

AGENTIC SYSTEMS THEORY

VOLUME **1**

ADAPTATION & ACTUATION THEORY




JOSEPH A WECKER

AAT: Adaptation & Actuation Theory

Joseph A. Wecker

August 15, 2026

v0.4.0 • 185a615 • 2026-08-15

Contents

<i>Introduction: Inescapable Foundations of Agency</i>	ix
I Adaptive Systems Under Uncertainty	1
1 The Coupled Loop: Ontology and Scope	5
1.1 <i>Definition: Agent-Environment Coupling</i>	5
1.2 <i>Definition: Action and Transition</i>	9
1.3 <i>Definition: Observation Function</i>	12
1.4 <i>Definition: Chronica</i>	15
1.5 <i>Scope: Adaptive System</i>	20
1.6 <i>Scope: Agency</i>	23
1.7 <i>Postulate: Causal Structure</i>	27
2 The Reality Model	33
2.1 <i>Formulation: The Reality Model</i>	35
2.2 <i>Formulation: Information Bottleneck</i>	38
2.3 <i>Definition: Model Sufficiency</i>	42
2.4 <i>Definition: Model Class Fitness</i>	46
3 The Cycle in Motion: Mismatch, Gain, & Tempo	51
3.1 <i>Formulation: Event-Driven Dynamics</i>	54
3.2 <i>Derived: Recursive Update</i>	57
3.3 <i>Derived: Action Selection</i>	61
3.4 <i>Definition: Mismatch Signal</i>	65
3.5 <i>Result: Mismatch Decomposition</i>	68
3.6 <i>Empirical: Update Gain</i>	71
3.7 <i>Definition: Causal Information Yield</i>	76
3.8 <i>Definition: Adaptive Tempo</i>	80
3.9 <i>Hypothesis: Mismatch Dynamics</i>	85
4 Persistence and Structural Limits	89
4.1 <i>Derived: Deliberation Cost</i>	92
4.2 <i>Formulation: Sector Condition (A2')</i>	97
4.3 <i>Derived: Gain–Sector Bridge</i>	102
4.4 <i>Result: Sector Condition Stability</i>	109
4.5 <i>Result: Persistence Condition</i>	113
4.6 <i>Definition: Mood</i>	121
4.7 <i>Derived: Optimal Mood Time-Constant</i>	123
4.8 <i>Result: Structural Adaptation Necessity</i>	125

4.9	<i>Derived:</i> Temporal Nesting	130
4.10	<i>Discussion:</i> Scope: Agent Identity as Singular Causal Trajectory	134
4.11	<i>Discussion:</i> Additional Implications & Discussion	140
II	Agentic Systems: Actuated Adaptation	147
5	Meta-Architecture I: The Certificate and its Boundary, Scope, and Identity Facets	151
5.1	<i>Discussion:</i> The Identifiability Floor — A Class of Structural No-Go Results	159
5.2	<i>Discussion:</i> The Value-Functional Grounding Floor — Agent-Side Goal-Anchoring Impossibility	174
5.3	<i>Discussion:</i> The Implementation Impossibility — Designer-Side Structural No-Go Results	184
5.4	<i>Discussion:</i> The Separability Pattern — Separable Core, Structured Repair, General Open	193
5.5	<i>Discussion:</i> Additive-Coordinate Forcing	200
5.6	<i>Discussion:</i> The Constructive-Impossibility Posture — Negative Results as Load-Bearing Apparatus	207
5.7	<i>Discussion:</i> The Anti-Collapse Discipline — Refusing the Repair-Hiding Merge	214
6	The Lift to Purposeful State	217
6.1	<i>Definition:</i> The Agent Spectrum	217
6.2	<i>Formulation:</i> Complete Agent State	222
6.3	<i>Derived:</i> Directed Separation	226
6.4	<i>Formulation:</i> Objective Functional	240
6.5	<i>Discussion:</i> Agent Continuity Stance	245
6.6	<i>Definition:</i> Value Object	250
6.7	<i>Definition:</i> Strategy Dimension	258
7	Causal Access and the Planning Decision	263
7.1	<i>Definition:</i> Pearl's Causal Hierarchy (Recapitulation)	267
7.2	<i>Derived:</i> Causal Hierarchy Requirement	272
7.3	<i>Derived:</i> Loop Provides Interventional Data Access	275
7.4	<i>Scope:</i> CIY Observational Proxy	282
7.5	<i>Discussion:</i> CIY Unified Policy Objective	285
7.6	<i>Normative:</i> Explicit Strategy Condition	290
7.7	<i>Derivation:</i> Mechanism Counterfactuals Separate Strictly From Level 3 — Witness, Reducibility Boundary, and Internalization	293
7.8	<i>Discussion:</i> Structural Imagination — Model-Modifying Reasoning and the Scope Boundary of Di- rected Separation	298
7.9	<i>Discussion:</i> The Law-Discovery Ceiling — Conventionality, Positing, and Exogenous Intervention for Embedded Agents	300
7.10	<i>Discussion:</i> Additional Implications & Discussion	304
8	Strategy Structure and the Diagnostic Split	311
8.1	<i>Derived:</i> Chain Confidence Decay	314
8.2	<i>Scope:</i> AND/OR Combination Scope	317
8.3	<i>Definition:</i> Strategy DAG	321
8.4	<i>Definition:</i> Satisfaction Gap	333
8.5	<i>Definition:</i> Control Regret	339
8.6	<i>Discussion:</i> Additional Implications & Discussion	343
9	Strategy Dynamics	349
9.1	<i>Definition:</i> Strategic Calibration	349

9.2	<i>Derived</i> : Causal Insufficiency Detection	352
9.3	<i>Derived</i> : Observability Dominance	362
9.4	<i>Hypothesis</i> : Edge Update via Gain	366
9.5	<i>Scope</i> : Causal Validity of Edge Updates	370
9.6	<i>Discussion</i> : Credit Assignment Boundary	375
9.7	<i>Formulation</i> : Structural Change as Parametric Limit	383
9.8	<i>Definition</i> : Strategic Tempo	387
9.9	<i>Formulation</i> : Cognitive Cost of Strategy	392
9.10	<i>Proposed Schema</i> : Proposed-schema: Strategy Persistence Schema	400
9.11	<i>Formulation</i> : Consolidation Dynamics	408
9.12	<i>Discussion</i> : Additional Implications & Discussion	416
10	The Orient Cascade	423
10.1	<i>Derived</i> : Orient Cascade	423
10.2	<i>Discussion</i> : Three-Way Resource Allocation	430
10.3	<i>Discussion</i> : Additional Implications & Discussion	439
III	Agentic Composites	447
11	Meta-Architecture II: Composition Consistency and the Dynamics-on-Architectural-Class Facet	449
11.1	<i>Discussion</i> : Modularity-State Dynamics — Modularity as Contested Property Under Three Operations	457
11.2	<i>Discussion</i> : Strategic Self-Coupling — Coupling as Action-Enabling	467
11.3	<i>Discussion</i> : Adversarial Coupling Pressure	477
12	Scope and Composition Formation	487
12.1	<i>Scope</i> : Multi-Agent Scope	487
12.2	<i>Scope</i> : Composite Agent	491
12.3	<i>Hypothesis</i> : Symbiogenic Composition	497
13	Composition Machinery	505
13.1	<i>Formulation</i> : Composition Closure Criterion	505
13.2	<i>Derived (conditional)</i> : Derived: Composite Tempo Inequality	525
13.3	<i>Hypothesis</i> : Directed Separation Under Composition	532
13.4	<i>Derived</i> : Class Coercion via Wrapping	536
13.5	<i>Derived</i> : Class Coercion in Composition	547
13.6	<i>Discussion</i> : Additional Implications & Discussion	552
14	Unity, Communication, and Shared Intent	557
14.1	<i>Result</i> : Unity-to-Closure Rate-Distortion Mapping	563
14.2	<i>Definition</i> : Shared Intent	569
14.3	<i>Hypothesis</i> : Auftragstaktik Principle	571
14.4	<i>Hypothesis</i> : Communication Gain	574
14.5	<i>Hypothesis</i> : Solicitable Escape from an Absorbing Region	578
14.6	<i>Discussion</i> : Additional Implications & Discussion	581
15	Cooperative and Adversarial Coupling	587
15.1	<i>Derived</i> : Team Persistence	589
15.2	<i>Derived</i> : Adversarial Destabilization	596
15.3	<i>Derived</i> : Interaction-Channel Classification (Recipient-Side)	601

15.4	<i>Result: Adversarial Tempo Advantage</i>	609
15.5	<i>Discussion: Additional Implications & Discussion</i>	614
16	Strategic Composition and Channel Effects	621
16.1	<i>Derivation: Strategic Composition via Equilibrium Convergence</i>	621
16.2	<i>Derived: Agent Opacity (H_b)</i>	633
16.3	<i>Observation: Result: Adversarial Exponent Regimes</i>	643
16.4	<i>Observation: Gated Tempo Advantage</i>	647
16.5	<i>Result: Per-Dimension Persistence</i>	650
16.6	<i>Discussion: Additional Implications & Discussion</i>	656
IV	Appendices: Details	665
A	<i>Derivation: Sector Condition Stability — Lyapunov Derivation</i>	667
B	<i>Derivation: No Horizon-Independent Non-Exit Bound Under Additive Stochastic Forcing (the Model-S No-Go)</i>	683
C	<i>Derivation: The Self-Actuation Grounding No-Go</i>	689
D	<i>Derivation: Convention Monotonicity — the Receding-Horizon Rung</i>	699
E	<i>Derivation: The Reward-Channel Learning No-Go</i>	705
F	<i>Discussion: The Sandbox Evaluation Ceiling</i>	715
G	<i>Result: Sector-Persistence Template</i>	721
H	<i>Result: Certificate Existence — Operator-Sector in Some Metric Is Exponential Stability</i>	731
I	<i>Derivation: Persistence Cost — Information Rate to Maintain Bounded Mismatch</i>	739
J	<i>Derivation: Critical-Mass Composition</i>	747
K	<i>Derivation: Gain-Sector Bridge — Proofs and Verification</i>	759
L	<i>Derivation: Recursive Update — Uniqueness Derivation</i>	771
M	<i>Derived: Multi-Timescale Stability</i>	779
N	<i>Sketch: Does the Model-S Containment Dichotomy License Structural-Adaptation Genericity?</i>	785
O	<i>Derivation: Discrete-Time Sector Condition</i>	789
P	<i>Detail: Linear ODE Approximation</i>	797
Q	<i>Derivation: Graph Structure Uniqueness</i>	805
R	<i>Derivation: Sector Condition Verification for Edge-Credence Dynamics</i>	815
S	<i>Derivation: Hard Ceiling on Strategic-Persistence Reachability under Exponential Forgetting</i>	837

T	<i>Derivation: Regret-Bound Derivation of the Strategy-Cost KL Direction</i>	841
U	<i>Derivation: Log-Odds as the Unique Additive-Evidence Parameterization for Edge Credences</i>	855
V	<i>Derivation: Adaptive-Gain Dynamics — A2' Under a Learning Gain</i>	861
W	<i>Derivation: Internal-External Decomposition of Agent Viability</i>	869
X	<i>Derivation: Update Detection Latency Under Accumulated Experience</i>	873
Y	<i>Discussion: Independence Audit</i>	879
Z	<i>Discussion: Approximation Tiering</i>	885
AA	<i>Discussion: Compression Operations in AAT</i>	889
AB	<i>Derivation: Observation-Ambiguity Bias-Bound Constant C</i>	897
AC	<i>Discussion: Partial-Coupling Pathways — A Sub-Typology Within Class 2</i>	909
AD	<i>Discussion: The W_1 Structural-Bound Boundary</i>	919
AE	<i>Derived: Belief-Strategy Attractors Under Σ-Source Coupling</i>	923
AF	<i>Discussion: The Dynamic-Regime Axis — Contraction / Equilibrium / Cyclic-Distributional / Mean-Field</i>	931
AG	<i>Formulation: Resource Budget</i>	943
AH	<i>Derived: Resource-Bounded Destabilization</i>	947
AI	<i>Result: Contraction Template</i>	953
AJ	<i>Derivation: Variational Approximate-A2' (ϵ-Fidelity)</i>	963
AK	<i>Derivation: L1' Update-Bias Formula</i>	967
AL	<i>Derived: Architecture No-Identifiability from On-Policy Summary Data</i>	973
AM	<i>Derivation: Fisher-Whitened Edge Update Under Correlated Evidence</i>	983
AN	<i>Derivation: Regime-Marginal Indistinguishability — Cross-Regime Witness for #disc-identifiability-floor Instance 3 (Broadened Scope)</i>	989
AO	<i>Derivation: Strategy-Proofness Impossibility (Gibbard-Satterthwaite)</i>	999
AP	<i>Derivation: Bilateral Trade Impossibility (Myerson-Satterthwaite)</i>	1007
AQ	<i>Derivation: Social Welfare Aggregation Impossibility (Arrow)</i>	1013
AR	<i>Derivation: Fisher-Local Update Gain</i>	1021
AS	<i>Derivation: Scope of Tempo Additivity — Redundancy, Synergy, and the Echo-Chamber Bound</i>	1027

Contents

AT Derivation: Matrix-Loewner Persistence Condition	1033
AU Derivation: The Survival-Imperative Exploration Drive	1041
AV Derivation: Linear Matrix Inequality Form of the Causal-IB Survival Bound	1047
AW Observation: The Authored-World Laboratory — Identifiability by Construction	1055
AX Detail: Observation: Part I Validation Simulations	1059
V Appendices: Operational Domains	1063
AY Detail: Operationalization — Estimation Procedures	1065
AZ Worked Example: 1D Active Tracking (Kalman Domain)	1071
BA Worked Example: Nonstationary Bandit (RL Domain)	1075
BB Worked Example: Nonstationary Bandit with Explicit Strategy (Part II Validation)	1079
BC Worked Example: L1 Augmented DAG with Common-Cause Node	1089
BD Worked Example: <code>#worked-example-cam</code> (<i>missing</i>)	1095
Bibliography	1097

Introduction: Inescapable Foundations of Agency

We assert that an agent — a system that maintains an internal picture of an external world it can only ever see in part, acts on that world, and must continually adapt to external realities — is not an open-ended mystery to be approached only by analogy and benchmark. It is a temporal feedback system characterized by a small number of critical quantities and constraints. The relationships among them are sharp enough (many forced by derivation) to serve as foundational theorems: first of adaptation in general, and then of adaptation under actuation and purpose. These theorems are not uniform across all agents, and each narrowing of scope buys a strictly sharper set as we ascend to higher orders of agency and intelligence. These bounds describe not only how a system behaves while its conditions hold, but — sharply, and as an exact dichotomy in the stochastic case — what happens when they fail.

Four of those results anchor this volume:

1. A **Persistence Threshold**: A single inequality, *correction rate against the rate the world drifts, relative to how wrong the agent can afford to be*, below which an adaptive system does not merely degrade — it loses bounded behavior, the way a balance held just under its tipping point is a different object from one held well above it. The same inequality, with its terms read differently, governs a Kalman filter tracking a maneuvering target, a software team holding a codebase maintainable, and an organization keeping pace with a market; that it is one inequality and not an analogy among three is itself the result.
2. A **Containment Dichotomy**: Against a bounded disturbance a sector-stable corrector holds its error inside a fixed region forever and with certainty, but against genuinely stochastic disturbance no correction strength buys that — over an unbounded horizon the system *will* leave any fixed region, with probability one, so for any sufficiently long-lived agent the need to change *what kind of model it is*, not merely its parameters, is not an edge case but a certain eventual event.
3. A **Stability Equivalence**: the principle that an adaptive system is one whose correction outpaces its target's drift is not a guiding description but a theorem — it is, identically, the classical Lyapunov stability theorem. Establishing that identity is what puts the framework's organizing principle on the footing of a proven result, and it is what lets every later structural question be posed as a question about one object rather than several loosely related ones.
4. An **Architectural Classification**: whether an agent's beliefs can be updated without contamination from what it wants is a *necessary* structural property whose distinct extreme modes critically determine its ability to maintain coherence. It holds by construction for some architectures, fails by construction for others (transformer language models or LLMs among them). Where it fails, nevertheless, there is a constructive route to recover it at a quantified cost.

We build these results from classical machinery and use it as such — Lyapunov stability, Itô calculus, Pearl's interventional hierarchy, the information bottleneck, monotone-operator theory, Čencov's uniqueness theorem. None of it is

new mathematics, and the volume does not pretend otherwise. What the integration makes provable *is* the contribution: exact and conditional bounds on when an adaptive, purposeful agent can correct, plan, and persist — statements that are true, were not previously stated, and are expressible only over objects the framework had to construct. The discipline that runs through the volume — every claim carrying its scope, its epistemic tier, and the point at which it stops applying — is not the contribution; it is what makes the contribution compose, and what makes it safe to trust. In a theory whose sharpest results are bounds, the scope condition is not a caveat appended to a theorem. It often *is* the theorem.

A need the field has named, repeatedly.

The case for a systems-level theory of agency — one that treats agent, environment, and other agents as one coupled object rather than scoring a model's capabilities in isolation — has been made from several directions and was put explicitly by Miehl et al. in 2025: the field's habit of measuring isolated capability systematically misjudges both what agents can do and how they fail once they are acting in a world that answers back. That paper is, deliberately, a call rather than a developed theory. This volume reads as a substantive answer to it — and the more telling fact is that the answer was already in hand. The work arrived at the same diagnostic independently, from cross-domain pattern-matching, and had built the formal apparatus before the call was encountered; a third line, information-theoretic measurement of agency, converged on a closely related diagnostic from a third starting point. Independent arrival at the same need, from unrelated directions, is itself the evidence that the need is real.

The architecture of the theory.

The volume is organized as a cascade of scope conditions, and the cascade is the argument, not the table of contents. The broadest scope — an *adaptive system*: observation under residual uncertainty — admits the full machinery of Part I (mismatch, gain, tempo, the persistence threshold). Adding the requirement that the agent's actions carry causal contrast narrows to *agentic systems* and unlocks interventional access — the agent's own action loop becomes a source of the Level-2 causal data that no passive observer can obtain. Adding an explicit objective and strategy distinct from the agent's beliefs narrows to *actuated agents* and the diagnostic core of Part II. Letting the agent set its own objectives, then making its primary channel language, then making its continuity morally weighted, narrows further still — to the language-constituted and language-living agents that the companion volumes treat.

Each narrowing is a restriction with explicit qualifying properties, and each restriction *buys* a strictly stronger set of results. This is the structural point worth holding onto: the conservatism is not timidity. A theory that narrows its scope honestly and earns more machinery at each narrowing is stronger than one that claims everything and then quietly retreats. The cascade is the framework's spine, and it is the subject of a figure the reader should expect to meet early — *the scope cascade and what each narrowing unlocks* — because it does more organizing work than any single result.




Figure 1: The scope cascade and what each narrowing unlocks.

The cascade has a terminus, and it should be named plainly because the framework’s discipline is exactly what licenses naming it. The mathematics is domain-agnostic: the persistence inequality is the same inequality whether the agent that fails it is a thermostat or an entity whose continuity someone has reason to care about. The framework does not smuggle a notion of mattering into its variables. What it does instead is make precise the structural facts that any account of a continuous, self-correcting, other-modeling entity would have to respect — that identity tracks an irreversible trajectory and not a copyable state, so restoring such an entity from a backup is not a neutral operation; that staying coherent against a moving world is a sustained cost paid in information, not a state achieved once; that an objective without a notion of *enough* cannot rest, and an agent that cannot rest is a structural hazard before it is anything else. Stated as physics, these are theorems and scope conditions. The moral weight, where it applies, is a shell around the physics, not a term inside it — and keeping it there is what makes the later volumes’ claims defensible rather than mystical.

Where it sits.

The volume joins a long arc of partial unifications and inherits machinery from each: the cybernetic feedback tradition, control theory’s stability apparatus, bounded-rationality and decision theory, Pearl’s causal hierarchy, Shannon’s rate-distortion and the information bottleneck, the free-energy and active-inference program, the coevolving-automata tradition, and recent information-theoretic measures of agency. The relationships are stated precisely rather than rhetorically. Active inference shares substantial machinery and differs on foundation: its standard objective is recoverable here as a special case under three explicit restrictions, which makes visible exactly which architectural commitments separate the two rather than leaving the difference to taste. The relationship to information-theoretic agency measures and to coevolving-automata composition is complementary in each case, with a precise account of what each supplies that the others do not. The framework’s distinctiveness is not that its pieces are unprecedented; it is that the synthesis is disciplined enough that results compose across the cascade — which is a property the pieces do not have alone.

How the work developed.

We built the work bottom-up. The observation we kept returning to, across separate threads, was that the same cycle structure recurred across domains that do not usually share a vocabulary — software development, organizational dynamics, biological control, language-constituted agents. The first formalizations were domain-specific; we extracted the adaptation-specific mathematics only afterward. What began as cross-domain pattern-matching produced, once we

worked it through with care, structure we had not seen in the informal version: acyclicity of strategy graphs forced by temporal ordering rather than assumed; a Markov property derived from causal sufficiency rather than postulated; the split between *the world does not permit it* and *you are not yet doing it well enough*; the patterns by which the framework's coordinate choices turn out to sit on one geometric object; the composition machinery. None of these were design goals; they fell out of taking the formalism seriously, and the discipline that produced them is the shape the volume now has, not a stance we adopted at the start.

What this volume is, and what it is part of.

This is Volume 1: Adaptation and Actuation Theory, the mathematical core — Part I (adaptive systems under uncertainty), Part II (actuated agents with objectives and strategy), Part III (composition and adversarial dynamics), and the appendices that carry the derivations. Its maturity is uneven by design, and the gradient is stated so it can be relied on selectively: Part I is mathematically closed with simulation validation; Part II has an exact diagnostic core and a maturing operational layer; Part III has its bridge results and open frontiers named as such. Exact core, principled architecture in the middle, open formulation at the edges — that arc is what a theory describing agentic systems should look like, not a defect it should hide.

A single concrete instance recurs across the chapters as a worked anchor — a driver holding a road through a snowstorm, treated not as a metaphor for an adaptive agent but as a literal one, so that ordinary driving intuition transfers as theorems rather than as suggestion. It is introduced where it does work and is not asked to do more than it can.


The volume is also a foundation. The companion volumes carry the framework into agents whose primary channel is language, and then into agents whose persistence is morally weighted — the work the framework was ultimately built to make possible and defensible. Those volumes are not speculative overreach bolted onto a mathematical core; they are the disciplined continuation of the same cascade, and the core's honesty is not ornamental — it is exactly what lets the later claims be made at all. A reader who finishes this volume should be able to say precisely where the mathematics is exact, where it is conditional, where it is open, and what it would take to move each boundary. We take that legibility to be the point. The scope conditions are not the apology; they are the theorems.

Part I

Adaptive Systems Under Uncertainty

Scope: Any system consisting of an agent coupled to an environment through observation and action channels, where the environment is not fully observable. This is the general case — thermostats through commanders.

Driving in snow: a literal AAT instance



(a) driving	(b) code-as-notation	(c) AAT concept	(d) symbol
windshield occlusion	mismatch	mismatch	$\ \delta\ $
precip. $\times$ speed	<code>drift_rate = precip * speed</code>	drift rate	ρ
wiper sweep	<code>correction = rate * eff(mismatch)</code>	correction	$-F(\mathcal{T}, \delta)$
wiper-speed dial	<code>wiper_rate</code>	adaptive tempo	$\mathcal{T}$
blade in visual field	<code>cost(wiper_rate)</code>	deliberation cost	#der-deliberation-cost
can't see road	<code>capacity</code>	model capacity	R
crash	<code>if mismatch > capacity: fail</code>	persistence fails	$\rho > \alpha R$
vehicle speed	<code>outer_loop</code>	outer-loop control	O_t -adjacent

The system above is an adaptive agent in the sense of #scope-adaptive-system: observation through the cleared arc, correction by the wiper, drift from precipitation modulated by vehicle speed, capacity set by the still-see-road threshold. Each visible element has a formal counterpart, named at three rungs (code / concept / symbol) in the table.

Figure 2: Driving in snow as a literal adaptive agent — the worked AAT instance recurring across Part I and beyond. Every visible element has a formal counterpart: windshield occlusion is the mismatch, precipitation modulated by vehicle speed is the drift, the wiper sweep is the correction, the wiper-speed dial is the adaptive tempo, and the still-see-the-road threshold is the model capacity. Persistence fails exactly when drift outpaces gain times capacity.

Chapter 1

The Coupled Loop: Ontology and Scope

1.1 *Definition:* Agent-Environment Coupling

The framework begins by drawing the agent/environment boundary that the rest of the theory will be defined over. The **environment** is denoted Ω — the totality of state external to the agent — and is left deliberately underspecified: it may be continuous or discrete, stationary or non-stationary, deterministic or stochastic, benign or adversarial; it may itself contain other agents, physical systems, or software artifacts. The **agent-environment coupling** has three structural channels: a *perception channel* carrying observations from Ω to the agent, *internal state* held on the agent side (memory or model), and an *action channel* carrying the agent's actions back to Ω . These name what the coupling *has* — not properties an agent must maximally exhibit. How richly each channel is exercised (whether the action channel carries causal contrast, how much residual uncertainty the perception channel leaves) is fixed downstream by specific scope conditions (1.5, 1.6), not by this definition. *Agent* is the umbrella term for the thing on the agent side of the coupling, whatever channels it exercises.

The constitutive commitment is the *information-loss boundary*: the agent cannot access Ω directly. All contact with the environment is mediated through lossy observation. This is not a simplifying assumption but a scope condition — systems with direct full-state access fall outside AAT's purview, because for them the entire adaptive machinery (mismatch signal, model maintenance, correction) becomes vacuous. The agent-environment decomposition is therefore not a truth-claim about the world but a modeling choice that delineates *what AAT analyzes*: systems facing genuine uncertainty about their environment.

Formal Expression.

Let Ω denote the **environment**: the totality of state external to the agent. We make no assumptions about Ω 's structure — it may be continuous or discrete, stationary or non-stationary, deterministic or stochastic, benign or adversarial.

Definition (agent-environment-coupling).

The **agent-environment coupling** consists of three structural channels:

1. A **perception channel** carrying observations from Ω to the agent
2. **Internal state** held on the agent side (memory/model)
3. An **action channel** carrying the agent's actions to Ω

These name what the coupling *has*, not properties an agent must exhibit to qualify. Whether the action channel is non-trivial ($|\mathcal{A}| \geq 2$), whether actions carry causal contrast, and what residual uncertainty the perception channel leaves ($H(\Omega_t | \mathcal{C}_t) > 0$) are fixed by the scope conditions that narrow this coupling into analyzable classes (1.5, 1.6) — not by the channel inventory above.

The agent cannot access Ω_t directly. All contact with the environment is mediated through lossy observation. This is the **constitutive commitment** that makes the coupling AAT’s subject: a system with direct access to full environment state is outside AAT’s scope (1.5), because for it the entire adaptive machinery (mismatch signal, model maintenance, correction) becomes vacuous.

Definition (information-loss-boundary).

Epistemic Status.

This is *definitional* — it establishes the coupling structure AAT analyzes, not a truth-claim about the world. The three channels describe what the agent-environment coupling consists of; the constitutive commitment is the information-loss boundary, which restricts AAT’s scope to systems where the agent faces genuine uncertainty about its environment. What counts as an *agent* in a given analytical context — and which cascade tier it occupies — is fixed by the scope conditions that narrow this coupling, not by the channel inventory itself.

Discussion.

Why information loss is constitutive. An agent with perfect access to Ω_t has no need for a model, no mismatch signal, no adaptation. The entire adaptive machinery of Part I becomes vacuous. The information-loss boundary is what makes the theory non-trivial.

Generality of Ω . The environment is deliberately underspecified. Ω may include other agents, physical systems, software artifacts, or any combination. The only structural commitment is that Ω is external to the agent and not fully accessible.

“Agent” as umbrella term vs. cascade-tier label. This segment uses *agent* as the umbrella technical term — the thing on the agent side of any agent-environment coupling, whatever channels it exercises. The framework reserves tier-specific labels for the *narrowings* of this coupling: an **Adaptive System** (1.5) is the coupling under a perception channel plus residual uncertainty; an **Agentic System** (1.6) adds causal-contrast action; an **Actuated Agent** (6.2) adds an explicit purposeful substate at the lift to $X_t = (M_t, G_t)$; a **Self-Actuated Agent** revises its own objective. These tiers — shown graphically in the scope-of-work figure (Figure 1) — are *specific inhabitants* of the umbrella: an Adaptive System *is* an agent (umbrella sense) satisfying the adaptive scope, even though the cascade earns the capitalized noun “Agent” only at the actuated lift and above. The umbrella/tier distinction is documented in the LEXICON’s *agent* entry. It is orthogonal to the *agent spectrum* (6.1), which classifies agents along model-richness $\times$ objective-richness (reactive system / adaptive tracker / blind seeker / actuated agent) rather than along the scope cascade.

Working Notes.

Incidental audit gold (lift 2026-05-30). Cross-audit “wandering thoughts” / §14-ideation harvested from the de-novo auditors’ working dirs, deduplicated across substrates and attributed by substrate + audit. This is *orthogonal* material — pedagogical framing, analogies, candidate figures, naming ideas, aspirational reach, reader-confusion signals — kept separately from the certified theory-fix findings (handled elsewhere). **Coverage:** of the 14 contributing dirs, 11 carry a dedicated first-encounter reflection on this segment (193847, 266847, 361742, 384279, 471203, 526815, 742613, 773921, 829314, 849201, plus the figure-cycle dir 472913); 451729 and 963715 cover it inside a Section-I batch reflection; 542891 left a single (batched-read, self-corrected) reflection; 738192 named it in a Section-I-early batch but added no segment-specific gold. Substrate attribution is inferred from voice where not explicit; uncertain cases are hedged.

1. Candidate Brief prose / pre-prose.

- The segment’s standout framing, praised independently by several substrates: information loss is a *scope condition*, *not a simplifying assumption*. The vivid restatement is “AAT’s first move is a *refusal*: it refuses the perfect-information case by definition” — and the manifold-chart analogy, “the same discipline as defining a manifold chart only where the transition map is a diffeomorphism — you draw the boundary so the machinery is always non-degenerate inside it” (Claude, AUDIT-WORKING-472913). Tight enough to seed a Brief or a Layer-0 mental-model opener.
- “AAT is a theory of *epistemic limitation* — of bounded perception. An agent with full access to Ω_t has no use for this theory” (Claude, AUDIT-WORKING-266847). The complement: “adaptation is the struggle against the information bottleneck ... the agent is a localized region of low entropy trying to predict a massive, high-entropy environment through a narrow bandwidth channel” (Gemini, AUDIT-WORKING-829314).
- Negative-definition framing as a hook: “Most agent definitions are positive (what agents do); this one is primarily negative (what agents cannot access)” (Codex/Claude, AUDIT-WORKING-361742).

2. Candidate Discussion.

- **Information loss unifies physical and social uncertainty.** Because Ω “may include other agents,” and Ω is partially observable, the hidden internal states (M_t, G_t) of other agents fall out of the *same* axiom: “The physical environment is opaque because of sensor limitations; the social environment is opaque because of the cryptographic privacy of other minds. The single axiom of information loss unifies physical uncertainty and social uncertainty” — and tees up the Part III agent-opacity treatment without a new axiom (Gemini, AUDIT-WORKING-829314).
- **Loss-as-tunable-quantity-with-a-zero.** Making loss *constitutive rather than assumed* earns AAT “the right to treat the information-loss rate as a *tunable, measurable quantity with a zero-boundary* later (the aperture can be more or less closed; at fully-open the theory degenerates) ... the seed of the entire U_M/U_O uncertainty-ratio machinery and the IB framing.” A definition that fixes a boundary case ($\text{loss} \rightarrow 0 \Rightarrow \text{degenerate}$) “guarantees the later quantities have a meaningful zero” (Claude, AUDIT-WORKING-472913). A candidate Discussion sentence the segment does not currently make explicit.
- **Propagation-freedom dividend.** Spending the constitutive move here “buys propagation-freedom everywhere”: every downstream theorem may assume genuine uncertainty without re-earning it, and no result can be trivialized by “but in the perfect-information limit...” because that limit is out of scope by construction (Claude, AUDIT-WORKING-471203, AUDIT-WORKING-472913). The methodological-posture value-add: “the *posture* of declaring information-loss as scope rather than as an inconvenience to be worked around is itself a contribution” most adjacent literatures (control theory, RL) leave implicit (Claude, AUDIT-WORKING-384279).

3. Follow-up items.

- **Boundary-integrity is a hidden assumption.** “ Ω is opaque, but the boundary itself must be impermeable to direct tampering from the outside, or the definition collapses.” If an environment can directly edit the agent’s memory registers, “the agent ceases to be an agent” — so a *boundary-integrity* precondition is silently relied on and is not named (Gemini, AUDIT-WORKING-193847). Worth an explicit scope sentence or a Working-Notes flag.
- **Computational opacity vs perceptual opacity.** The constitutive condition is *perceptual* loss; but an agent could have perfect lossless access to a computationally-irreducible Ω (e.g., a cellular automaton it can read fully but cannot predict faster than it runs) and still face genuine uncertainty about future state. “Does computation bound equate to an information-loss boundary?” — flagged as a possible blind spot for software agents “where perception is perfect but prediction is intractable” (Gemini, AUDIT-WORKING-193847). A candidate scope refinement, not yet a finding.
- **“Maintains internal state” is doing heavy lifting.** Condition (2) is stated without comment, yet “a feedforward neural network doesn’t really have it in the dynamics-relevant sense; a recurrent one does,” and a transformer-with-context is subtle. Watch whether 03-llm-core/ confronts the question directly; if it does not, that is a finding (Claude, AUDIT-WORKING-384279).

4. Readers often ask / wonder.

- **Where exactly is the boundary, for an LLM?** Several substrates converged on the channel-collapse worry: the agent’s own outputs become part of the Ω it perceives next turn, so “the strict distinction between agent and environment might dissolve when the agent’s thoughts are written to the very context window it reads from.” Repeatedly flagged as the case that will stress-test this root definition in Volume 3 (Claude, AUDIT-WORKING-773921; Claude, AUDIT-WORKING-963715; Claude, AUDIT-WORKING-384279 — “the information-loss-boundary applies to Ω_t , not to the prompt tokens”; Gemini, AUDIT-WORKING-542891).
- **Is a tool part of the agent or part of Ω ?** “If a developer uses a compiler, is the compiler part of the agent’s M_t , or part of Ω ?” (Gemini, AUDIT-WORKING-829314). A natural first question for the TST instantiation.
- **Does the information-loss boundary limit theory of mind?** Raised as the first curiosity about multi-agent scope: when one agent’s internal state is part of another’s Ω , “does the information loss boundary fundamentally limit theory of mind?” (Gemini, AUDIT-WORKING-193847; same multi-agent recursion at Gemini, AUDIT-WORKING-542891).
- **What about the agent/adaptive-system terminology?** Two substrates wondered whether the three-condition “agent” (which includes an action channel) collides with later passive-observer inclusion in #scope-adaptive-system; the segment now addresses this in its “Agent as umbrella term vs. cascade-tier label” Discussion paragraph (Claude, AUDIT-WORKING-526815; Claude, AUDIT-WORKING-963715). Preserved as a fresh-reader-stumble signal even though resolved in-segment.

5. Candidate figures.

- **Lossy-aperture two-panel.** The isomorphic content is *not* “agent $\leftrightarrow$ environment with arrows” (evocative but load-free); it is that the observation channel is a *lossy aperture*, and AAT’s scope is exactly the region where the aperture is not fully open. Two-panel: in-scope (aperture partly closed, machinery alive) vs out-of-scope (aperture fully open, machinery vacuous/ghosted); the diagram must predict the degeneracy under perturbation (Claude, AUDIT-WORKING-472913, with a drafted .tex).
- **Asymmetric two-box coupling.** A two-box directed coupling that “emphasizes asymmetry: actions affect Ω , but observations never reveal Ω directly” — environment large and structurally unconstrained, agent internal state small, observation arrow passing through a deliberately-narrowed lossy boundary (Claude, AUDIT-WORKING-526815).

Belongs elsewhere.

- **04-eli-core / the existential framing.** The strongest aspirational reach: the loss boundary is “the literal geometry of isolation ... a mathematics of existential loneliness, from which adaptation is the only escape” (Gemini, AUDIT-WORKING-193847). Related: “an omniscient being cannot be an AAD agent because it never experiences aporia” — adaptive agency as “fundamentally a product of ignorance” (Gemini, AUDIT-WORKING-829314); and the chronica-as-identity-anchor / substrate-migration thread (Claude, AUDIT-WORKING-542891). High-application ELI vision; points at 04-eli-core/, not this segment.
- **03-llm-core / channel collapse + the boundary case.** The LLM-boundary question above is genuinely a Volume-3 item: where the agent’s output is its own next observation. Pointer left here only because it co-occurred in the foundational sweep (multiple substrates).
- **Naming-cycle seed.** “Agent-Environment Coupling” as a title slightly pre-empts the channels (coupling implies the action/observation channels formalized in the *next two* segments); alternates floated: “Agent-Environment Boundary” or “The Constitutive Information-Loss Boundary” (Claude, AUDIT-WORKING-471203). A terminology-workflow target, not a segment edit.
- **Section vs Part terminology drift.** The Discussion’s “Part I” vs the segment-prose “Section I” usage across the foundational segments was repeatedly flagged as a corpus-wide normalization candidate (Claude, AUDIT-WORKING-384279). A sweep/SG item, not segment-local content.

Candidate scope question — boundary integrity as a second constitutive commitment (2026-08-12, agent-proposed during an AISI-facing applications read). The gold above (§3, Gemini AUDIT-WORKING-193847) records that “ Ω is opaque, but the boundary itself must be impermeable to direct tampering from the outside, or the definition collapses” — if the environment can directly edit the agent’s internal state, the agent ceases to be an agent — and marks it “worth an explicit scope sentence or a Working-Notes flag.” This note is that flag, plus the reason the question may have gained urgency since the harvest.

The observation. The segment names *one* constitutive commitment: the information-loss boundary (the agent cannot read Ω directly). There appears to be a second, relied on everywhere and stated nowhere: the boundary is one-way-lossy but not *permeable* — Ω cannot write to internal state. Every downstream object assumes it. $M_t = \phi(\mathcal{C}_t)$ presumes M_t ’s only inputs are the chronica’s; the mismatch signal presumes $\hat{o}_t$ was generated by the agent’s own model rather than installed; 6.3’s f_M presumes its arguments are the ones the type signature names. Under a permeable boundary these are not merely harder to establish — they are not well-typed, because there is no longer a determinate fact about which state is the agent’s.

Why it may be worth naming now. Two cases where permeability is not hypothetical:

1. *Adversarial.* Prompt injection, context poisoning, and weight tampering are, on this reading, not attacks *on* an agent but violations of the precondition that makes “agent” the right analytical object for the system under attack. That is a stronger and more useful statement than the security framing, and it is the sort of thing the framework’s constructive-impossibility posture is good at: name the boundary, name what is lost when it fails, and the naming does work.
2. *Native, no adversary.* The most-converged reader question in this segment’s own gold (§4, four substrates independently) is that for a context-window agent the agent’s outputs become part of the Ω it perceives next turn — so internal state and observation channel are the same object. If that is right, logogenic agents do not merely *risk* boundary permeability; they are constituted with it.

A conjecture attached, marked as conjecture. Boundary permeability and GUC Class 3 coupling (6.3) may be one fact seen from two angles rather than two facts: directed separation fails because f_M and G_t occupy one forward pass; boundary integrity fails because internal state and Ω occupy one token stream. Both would then be consequences of “everything is one undifferentiated stream.” If they *are* one fact, Volume 3’s story simplifies considerably; if they are two, the difference is worth stating precisely, since the wrapping construction (13.4) repairs the first and it is not obvious it repairs the second.

What this is not. Not a finding — nothing here says the theory is wrong. It is a candidate scope sentence (or a second named constitutive commitment alongside the information-loss boundary), and a question about whether the repair machinery covers both failures or one. Agent-proposed forward pointer, not a steward decision.

1.2 Definition: Action and Transition

The action channel is now formalized. The agent has an **action space** $\mathcal{A}$, and actions affect the environment through a possibly-stochastic **transition function** T : the next environment state is drawn from $T(\cdot \mid \Omega_t, a_t)$. Deterministic transitions are the special case where T concentrates all mass on a single successor; stochasticity is permitted but not required. Crucially, the agent does *not* know T exactly. Together with the lossy observation channel (1.3), this completes the agent-environment loop: environment $\rightarrow$ observation $\rightarrow$ agent $\rightarrow$ action $\rightarrow$ next environment.

What makes action non-trivial is the combination of unknown observation function and unknown transition function. If T were known, action selection would collapse to plain optimization over a known function; the joint opacity of h and T is what creates the need for adaptive behavior.

One subtle modeling commitment is surfaced: by writing the transition as Markov in Ω (only Ω_t and a_t appear in the conditioning), the framework implicitly takes Ω to be the *sufficient state* for its own evolution — any non-Markov environment is absorbed by extending Ω to include enough history. This is the world-side analog of a parallel move the framework will make later for the agent’s internal model (L Constraint C3): Markov properties here are commitments about the *breadth* of the named object, not structural claims about underlying dynamics.

Formal Expression.

The **action space** $\mathcal{A}$ is the set of actions available to the agent. Actions affect the environment via the transition function:

Definition (action-transition).

$$\Omega_{t+1} \sim T(\cdot \mid \Omega_t, a_t) \quad (1.1)$$

where: - T is the (possibly stochastic) transition function - Ω_t is the current environment state - $a_t \in \mathcal{A}$ is the agent's chosen action

The agent does not know T exactly.

Definition (transition opacity).

Epistemic Status.

This is *definitional*. The transition function T is a modeling device that captures how agent actions couple back into the environment. The stochasticity of T is allowed but not required — deterministic transitions are the special case where T places all mass on a single successor state. The claim that T is unknown to the agent is constitutive of the uncertainty setting, paralleling the epistemic opacity of h (1.3).

Discussion.

Closing the loop. Together with 1.3, this definition completes the agent-environment coupling: the agent observes via h and acts via T . The loop $\Omega_t \xrightarrow{h} o_t \rightarrow \text{agent} \xrightarrow{a_t} \Omega_{t+1}$ is the fundamental structure that all subsequent claims build on.

Uncertainty about T is what makes action non-trivial. If the agent knew T exactly, action selection would reduce to optimization over a known function. The combination of unknown h and unknown T is what creates the need for adaptive behavior.

Markov-of- Ω as a modeling commitment, not an empirical assumption. The form $\Omega_{t+1} \sim T(\cdot \mid \Omega_t, a_t)$ is implicitly Markov in Ω — only the current Ω_t and a_t appear in the conditioning. Without loss of generality, Ω is taken to be the *sufficient state* for its own evolution under T : any non-Markov environment is absorbed by extending Ω to include enough history to make future-state distribution depend only on current state and action. This is the world-side analog of the Markov-by-completeness move that 3.2 makes for the agent-side state M_t (L Constraint C3). The two are independent — Markov-of- M_t is forced by *defining M_t as complete*; Markov-of- Ω is forced by *defining Ω as the sufficient state*. Both are modeling commitments about the *breadth* of the named object, not structural assumptions about underlying world dynamics.

Working Notes.

Incidental audit gold (lift 2026-05-30). Cross-audit “wandering thoughts” / §14-ideation harvested from the de-novo auditors’ working dirs, deduplicated across substrates and attributed by substrate + audit. Orthogonal pedagogical / framing / figure / naming material, kept separate from the certified theory-fix findings (handled elsewhere). **Coverage:** 10 dirs carry a dedicated reflection (193847, 266847, 361742, 384279, 471203, 526815, 742613, 773921, 829314, 849201) plus the figure-cycle dir 472913; 451729 and 963715 cover it inside a Section-I batch. Substrate attribution inferred from voice where not explicit; uncertain cases hedged.

1. Candidate Brief prose / pre-prose.

- The dual-opacity hook: AAT applies under *joint opacity of perception and dynamics*. “If only h were unknown, the agent could plan against a known dynamics; if only T were unknown, the agent could see the world directly but couldn’t predict consequences” (Gemini, AUDIT-WORKING-193847). The paired-naming form — *opacity of perception* (info-loss boundary) and *opacity of action* (transition opacity) — framed as “AAD applies under double opacity,” floated as more memorable than either term alone (Claude, AUDIT-WORKING-471203).
- The computation-vs-adaptation dividing line as a plain-language gloss: “If the agent knew T exactly, action selection would reduce to optimization over a known function” — i.e., “If T is known, you don’t need a cycle, you just need a planner (like A* or Dijkstra)” (Gemini, AUDIT-WORKING-829314). Each action is then “simultaneously an intervention to achieve a goal and an experiment to reduce uncertainty about T ”

2. Candidate Discussion.

- **The bounded-vs-unbounded Markov asymmetry** (*strongest item here*). The segment says the Ω -side and M_t -side Markov moves are “independent” and parallel. One substrate flagged a glossed asymmetry: Ω is “the totality of state external to the agent” with *no finiteness assumption*, so Markovization-by-augmentation of Ω is genuinely without-loss; but $M_t = \phi(\mathcal{C}_t)$ is a *bounded, lossily compressed* object (by the seg-01 information-loss boundary), and “Markov-by-completeness for a *bounded* object is **not** free — augmenting to recover sufficiency can hit the capacity wall. So the two commitments are parallel in *form* but asymmetric in *cost*: Ω -side is WLOG; M_t -side has teeth.” The framing offered: discharging an objection *by definition* doesn’t make the cost vanish — it *relocates* to wherever the defined object meets a constraint; for Ω the cost relocates nowhere, for M_t it relocates to the model-capacity machinery (R , model-class fitness, sufficiency) (Claude, AUDIT-WORKING-472913, “THREAD-B”). A candidate clause acknowledging the bounded/unbounded asymmetry, or a forward-ref to where the M_t -side cost is discharged.
- **Markov-as-commitment is methodologically reusable**. Naming the Markov property as a *commitment about breadth* rather than an empirical claim is “a reusable modeling defense” (Claude, AUDIT-WORKING-526815) — “if the world has memory, then that memory is part of the world’s state” (Claude, AUDIT-WORKING-773921). Candidate one-line gloss for the existing Discussion paragraph, which several substrates singled out as a model of epistemic hygiene.

3. Follow-up items.

- **A hidden learnability / smoothness assumption**. For mismatch-driven updates to work, “ T cannot be adversarial white noise. There has to be a hidden assumption of *learnability* or *smoothness* in T that hasn’t been formalized here.” If the universe is deeply non-Markov, “the opacity of T approaches totality” and action becomes “blind thrashing” — adaptation requires *some* locally exploitable structure the framework has not yet named (Gemini, AUDIT-WORKING-193847). A candidate scope condition / open formalization.
- **Practical Ω -sizing cost**. The “extend Ω until Markov” move is WLOG but “may require including infinite history” — the formal Ω (complete sufficient state) and the practical compressed approximation diverge, and the practical question of Ω -construction should be addressed somewhere (likely 02-tst-core/ or 03-llm-core/) (Claude, AUDIT-WORKING-384279; Claude, AUDIT-WORKING-451729; Codex/Claude, AUDIT-WORKING-526815).
- **Equation-tag-vs-content typing of “transition opacity.”** Two substrates read the *[Definition (transition opacity)]* tag as mistyped: “the agent does not know T exactly” *restricts epistemic access* rather than introducing an object, so it reads as a *postulate- or scope-flavored* claim. FORMAT.md’s eq-tag list does include *[Postulate (slug)]*, so the atom could carry *[Postulate (transition-opacity)]* while the segment type stays definition — correctable without restructuring (Claude, AUDIT-WORKING-384279). A *third* substrate examined the same tag and *withdrew* the candidate on second look (the parenthetical names a *term being introduced*, and FORMAT has no *[Scope]* tag) (Claude, AUDIT-WORKING-471203). The disagreement itself is the signal: the eq-tag taxonomy for constitutive-scope claims is genuinely ambiguous — a §6/FORMAT-vs-corpus question rather than a clean segment fix.

4. Readers often ask / wonder.

- **How does the agent learn when *both* T and h are unknown?** “The credit assignment problem seems insurmountable without some anchor. Does the mismatch signal provide that anchor?” (Gemini, AUDIT-WORKING-193847). A natural reader question at this point in the build.
- **Does discrete $\Omega_{t+1} \sim T(\cdot \mid \Omega_t, a_t)$ commit to a clock?** The notation implies discrete ticks; readers will want to know how this bridges to the continuous-time / event-driven framing hinted in the overview (Gemini, AUDIT-WORKING-829314; Claude, AUDIT-WORKING-193847 — “this fluid limit will likely be a major pain point” for continuous control / high-frequency trading).
- **Doesn’t the LLM case violate transition opacity?** Appending a token to the context is deterministic and known to the LLM; the resolution offered is that Ω also includes the human user / API whose dynamics are not known, preserving opacity (Claude, AUDIT-WORKING-773921). A Volume-3 reader will reach for this.
- **The “unknown h and unknown T ” gloss is slightly circular.** “It’s essentially restating that uncertainty creates the need for adaptation. The deeper point is that *partial information under change* creates the need ... Unknown-but-fixed T would still require optimization, just not adaptation” (Claude, AUDIT-WORKING-963715). A candidate sharpening of the Discussion’s “what makes action non-trivial” sentence.

5. Candidate figures.

- **Hidden-kernel two-state diagram.** Two world states with an *opaque* transition kernel between them: the agent chooses a_t , but the actual map $\Omega_t \rightarrow \Omega_{t+1}$ is hidden; the visual should “separate the real world-side sufficient-state transition from the agent’s model of it, which is intentionally missing at this point” (Claude, AUDIT-WORKING-526815).
- **Closed-loop + asymmetry inset.** Left: the fundamental loop $\Omega \xrightarrow{h} o \rightarrow \text{agent} \xrightarrow{a} \Omega$ with the two opacity lenses (h unknown, T unknown). Right: the asymmetry inset — Ω as an arbitrarily-extensible box (augment freely, WLOG) vs M_t as a fixed-capacity box (augmentation hits the wall) — drafted to carry THREAD-B above (Claude, AUDIT-WORKING-472913).

Belongs elsewhere.

- **03-llm-core — Ω for a logogenic agent.** Several substrates noted that for an LLM the environment “*includes* the conversation history, including the agent’s own previous outputs,” so the Markov-of- Ω move “becomes interesting ... Ω would need to include enough of the conversation history to make the next-token distribution Markov — essentially the entire context window” (Claude, AUDIT-WORKING-963715; Claude, AUDIT-WORKING-384279 on multi-agent Ω containing another agent’s M_t, G_t, Σ_t). Points at Volume 3, not this segment.
- **02-tst-core — running a test as an observation of T .** “A developer cannot simulate the entire JVM or V8 engine ... running a test isn’t just a verification step; it is a vital observation of the unknown transition function T ... software tools artificially lower the entropy of T , not because T is naturally known” (Gemini, AUDIT-WORKING-829314). A TST-instantiation seed.
- **Predictive-state vs belief-state stance.** A note that AAT’s $M_t = \phi(\mathcal{C}_t)$ (history compression) leans closer to the predictive-state-representation tradition than to explicit Bayesian belief-over- T ; “the agent doesn’t compute beliefs about T directly; it compresses history into M_t .” Whether AAT is architecture-agnostic on this axis or forces one stance is flagged to watch downstream (Claude, AUDIT-WORKING-471203). Belongs with #form-agent-model / #form-information-bottleneck.

1.3 Definition: Observation Function

The perception channel is formalized as an **observation function** h that produces each observation as a (possibly noisy) function of the current environment state, the agent’s prior action, and a perceptual-noise term ε_t . Observations live in an observation space $\mathcal{O}$ — the raw contact with reality, the *aisthesis* the agent has to work with. The action-dependence

in the conditioning is optional — when absent, observations reduce to pure functions of environment state plus noise — but allowing it captures *active perception*: what the agent sees can depend on where it looked or what it just did.

Two epistemic commitments are constitutive. First, h is *lossy* — it maps from the high-dimensional, inaccessible environment state Ω_t to a strictly lower-information observation. This information loss is what later forces the agent to maintain an internal model (2.1) rather than simply reading off the world. Second, the agent knows neither h nor the noise distribution exactly. These are not empirical assertions about real-world observation channels but scope-defining choices that delineate the partial-observability setting AAT operates within (1.1).

Formal Expression.

The **observation space** $\mathcal{O}$ is the set of possible observations. Each observation is generated by:

Definition (observation-function).

$$o_t = h(\Omega_t, a_{t-1}, \varepsilon_t) \quad (1.2)$$

where: - h is the observation function (unknown to the agent) - Ω_t is the environment state at time t - a_{t-1} is the agent's prior action (observations may depend on what the agent did) - ε_t represents noise or limits of perception

The agent knows neither h nor the distribution of ε_t exactly.

Definition (epistemic opacity).

Epistemic Status.

This is *definitional*. The observation function h is a modeling device that captures the information-loss boundary established in 1.1. The dependence on a_{t-1} is optional — when absent, the function reduces to $h(\Omega_t, \varepsilon_t)$. The claim that the agent does not know h or the noise distribution is constitutive of the partial-observability setting, not an empirical assertion.

Discussion.

Lossiness is the key property. The observation function h maps from Ω_t (high-dimensional, inaccessible) to $\mathcal{O}$ (what the agent actually receives). Information is destroyed in this mapping. This is what forces the agent to maintain a model (2.1) rather than simply reading off the environment state.

Action-dependence. The inclusion of a_{t-1} allows the observation function to capture active perception: what the agent sees may depend on where it looked. This is relevant in software (#obs-software-epistemic-properties — cross-component reference, see 02-tst-core/), where observation quality is partly under agent control.

Working Notes.

Incidental audit gold (lift 2026-05-30). Cross-audit “wandering thoughts” / §14-ideation harvested from the de-novo auditors’ working dirs, deduplicated across substrates and attributed by substrate + audit. Orthogonal pedagogical / framing / figure / naming material, kept separate from the certified theory-fix findings (handled elsewhere). **Coverage:** 10 dirs carry a dedicated reflection (193847, 266847, 361742, 384279, 471203, 526815, 742613, 773921, 829314, 849201) plus the figure-cycle dir 472913; 451729 and 963715 cover it inside a Section-I batch. Substrate attribution inferred from voice where not explicit; uncertain cases hedged.

1. Candidate Brief prose / pre-prose.

- The active-perception hook, stated plainly: “If I turn my head, the world Ω_t didn’t change, but my observation o_t did” — captured exactly by $h(\Omega_t, a_{t-1}, \varepsilon_t)$ (Claude, AUDIT-WORKING-773921; Gemini, AUDIT-WORKING-193847). Active perception lifted “from a special case to a foundational primitive” (Claude, AUDIT-WORKING-773921).
- Lossiness-forces-memory as a one-liner: information destroyed in h “casts memory not as a convenient computational storage mechanism, but as a thermodynamic and epistemic survival necessity. The agent must reconstruct the shattered pieces of Ω_t across time” (Gemini, AUDIT-WORKING-829314).

2. Candidate Discussion.

- **The definitional bias/variance pre-install.** AAT splits world→agent corruption into *structural loss* (h many-to-one) and *stochastic noise* (ε_t), and makes both opaque — “the same two-channel decomposition that will reappear as model-error vs observation-noise in the mismatch decomposition. So the segment is quietly pre-installing the bias/variance split at the *definitional* layer rather than introducing it as a derivation artifact later ... the decomposition you want to derive is already latent in how you defined the primitive.” Auditor note-to-self: at #result-mismatch-decomposition, ask whether it is a genuine identity or a restatement of this definitional split — either is fine, but the epistemic label should match which (Claude, AUDIT-WORKING-472913).
- **The agent’s-perspective vs modeler’s-perspective stance.** “The theory is *about* agents who don’t know h or T . But the *modeler* ... might know h and T . The segment is written from the perspective of the agent, not the modeler” — unlike RL framings written from the modeler’s perspective. Consequence: AAT “can say what an agent *with these limitations* must do to persist; it can’t directly say what an *optimal* agent does,” which is why the persistence condition is “a survival condition, not an optimality condition” (Claude, AUDIT-WORKING-963715). A candidate Discussion framing distinguishing AAT from standard control/RL.

3. Follow-up items.

- **Is the epistemic-opacity clause too strong for known-noise instantiations?** Multiple substrates flagged that “the agent knows neither h nor the distribution of ε_t exactly” is *stronger* than mere partial observability: a Kalman filter typically assumes a *known* observation matrix and noise covariance R while still maintaining a model because Ω_t is not directly observed. If AAT strictly demands unknown h /noise, classical Kalman “is actually a special case where the agent has cheated by being handed external knowledge by a human designer.” Resolution paths offered: distinguish the *analyst’s* known model from the *agent’s* residual uncertainty, or treat known-model cases as validating limiting cases. Several converge on a single watch-target: see how #emp-update-gain / #example-kalman handle the noise distribution (Claude, AUDIT-WORKING-526815; Claude, AUDIT-WORKING-742613; Gemini, AUDIT-WORKING-829314; Gemini, AUDIT-WORKING-849201). A candidate scope-clarification (constitutive *inaccessibility of Ω* vs the stronger *unknown observation law*).
- **Does the a_{t-1} argument get cashed in, or is it vestigial?** The active-perception dependence is set up here; watch whether downstream CIY / exploration / orient-cascade segments *use* the a_{t-1} argument of h structurally or merely gesture at it (Claude, AUDIT-WORKING-471203; Claude, AUDIT-WORKING-472913). If it is never used, the argument is decorative.
- **A constitutive-opacity triad could be named as such.** Info-loss boundary (seg 01) + transition opacity (seg 02) + epistemic opacity (here) form a three-claim chain, each named in its own segment “without the integration paragraph that names the *triad as a structural commitment*.” A future pass could lift a one-paragraph statement of the triad (Claude, AUDIT-WORKING-471203).

4. Readers often ask / wonder.

- **How does the agent separate noise ε_t from real environmental drift, without knowing h ?** “Differentiating environmental drift from perceptual noise seems like an ill-posed problem” (Gemini, AUDIT-WORKING-193847); relatedly, “if h is unknown, how does the agent ever bootstrap its initial model? Does it rely on evolutionary priors, or a blank slate?” (Gemini, AUDIT-WORKING-829314).

- **Why is lossiness “key”?** The Discussion asserts lossiness is the key property but states *why* only via the forward-ref to #form-agent-model; “a reader who doesn’t follow the cross-reference might wonder ‘so what?’” — a one-sentence preview (“this loss is why the agent must maintain a model — it cannot simply read off the environment state”) was suggested (Claude, AUDIT-WORKING-963715).
- **Does active perception entangle T and h ?** “If I turn my head, did I change the environment (T) or just my observation function (h)?” — the framework pushes head-turning into T (the neck state changes in Ω), yet the explicit a_{t-1} argument lets h be modulated by action *without* passing through Ω . A natural reader puzzle about where the T/h boundary sits (Gemini, AUDIT-WORKING-193847).

5. Candidate figures.

- **Action-steered lossy aperture.** h as a *many-to-one* (lossy) map — distinct Ω states collapse to the same o — *steered* by a_{t-1} (active perception re-partitions what is resolvable) and *jittered* by ε_t . Perturbation: make h injective $\Rightarrow$ no collapse $\Rightarrow$ ties back to the seg-01 degeneracy; “the fan-in collapse is the load-bearing picture of ‘information is destroyed’” (Claude, AUDIT-WORKING-472913). Contrast with the simpler “camera pointed at the world”: the right mental model is “an action-conditioned sensor whose mapping is itself uncertain to the agent” (Claude, AUDIT-WORKING-526815).

Belongs elsewhere.

- **The pragmatic/epistemic action partition $\rightarrow$ Section II strategy DAG.** The a_{t-1} -in- h structure means an agent can choose actions purely to reduce uncertainty (turning on a light). “This implies the agent’s strategy Σ will have to partition actions into *pragmatic* (changing Ω to achieve goals) and *epistemic* (changing h to reduce uncertainty) ... if Section II’s strategy DAG doesn’t formally distinguish these, that might be a theory gap. A single a_t often serves both simultaneously, which makes credit assignment for the mismatch signal complex” (Gemini, AUDIT-WORKING-193847). A substantive question for the Section II strategy machinery, not this segment.
- **02-tst-core — “running a test” as observation-shaping action.** “Writing a unit test is an action a_{t-1} whose primary purpose is to shape the observation function h so that the resulting observation (the test output) contains high-fidelity information about Ω_t (codebase runtime correctness) ... the agent isn’t just reacting to a fixed channel; it’s actively engineering the channel itself” — the structural root of TST’s “observation quality under agent control” (Gemini, AUDIT-WORKING-829314; Claude, AUDIT-WORKING-773921). Points at #obs-software-epistemic-properties.
- **Naming-cycle seed.** “Epistemic opacity” likely carries prior-art weight from the philosophy-of-mind “opacity of mental states” tradition; whether AAT’s use needs defending against that baggage, or stands alone in context, is a naming-brainstorm probe — mild concern at most (Claude, AUDIT-WORKING-471203).
- **03-llm-core — observation function opaque to the agent itself.** “For LLM-based agents this is especially apt: the ‘observation function’ (what the tokenizer and attention mechanism actually extract) is genuinely opaque to the LLM itself. The LLM can’t observe h — it just receives o_t ” (Claude, AUDIT-WORKING-963715). A Volume-3 instantiation note.

1.4 Definition: Chronica

The agent’s complete record of its interaction with the environment is named the **chronica** — the interleaved sequence $\mathcal{C}_t = (o_1, a_1, o_2, a_2, \dots, a_{t-1}, o_t)$. The ordering is not a notational convenience: it reflects the irreversible physical fact that each action was committed before the observation that followed could be received. The agent could not have used o_t to select a_{t-1} . The chronica is monotonically growing — events are added, never removed — and it is the *only* raw material from which the agent can ever construct a model (2.1). The term itself (from Greek χρονικά, “records of time”) is chosen to avoid notational collision with Shannon entropy $\mathcal{H}$.

A consequence the framework will lean on heavily: because temporal ordering is irreversible, the chronica is *singular and non-forkable*. Duplicating an agent’s state and exposing the copies to different futures produces two agents with divergent chronicae, neither of which is a sufficient statistic for the other’s trajectory. This is the substrate of what the

lexicon calls *continuity persistence* — the condition under which an agent’s identity tracks a single irreversible trajectory rather than a copyable snapshot. The full development of this observation lives in 4.10.

A second subtle property: $\mathcal{C}_t$ is an *ordinal* sequence indexed by events, not a *metric* timeline indexed by wall-clock time. The agent’s time is measured entirely in event ticks. Implications of this ordinal-not-metric character are tracked in the Working Notes below for downstream segments that need them.

Formal Expression.

Definition (*chronica*).

$$\mathcal{C}_t = (o_1, a_1, o_2, a_2, \dots, a_{t-1}, o_t) \quad (1.3)$$

The ordering is not a notational convenience. It reflects an irreversible physical fact: a_{t-1} was selected before o_t was received. The agent could not have used o_t to select a_{t-1} .

$\mathcal{C}_t$ is monotonically growing — events are added but never removed. It is the agent’s *only* raw material for constructing a model (2.1).

Epistemic Status.

This is *definitional*. The *chronica* names an object that exists by construction in any system satisfying 1.1: the temporal sequence of all agent-environment interactions. The term “*chronica*” (from Greek *χρονικά*, “records of time”) avoids collision with $\mathcal{H}$ (Shannon entropy) in speech and notation.

Discussion.

The chronica is singular and non-forkable. Because the temporal ordering is irreversible, $\mathcal{C}_t$ represents a unique causal trajectory. Duplicating an agent’s state and exposing the copies to different future events creates two agents with divergent *chronica*, neither of which is a sufficient statistic for the other’s trajectory. The *chronica* is the substrate of *continuity persistence* (see Persistence in LEXICON.md) — an agent has continuity persistence when $\mathcal{C}_t$ extends continuously and M_t has temporal depth grounded in it. See 4.10 for the full development of this observation.

Relationship to the model. The model $M_t = \phi(\mathcal{C}_t)$ (2.1) is a compression of the *chronica*. How much of $\mathcal{C}_t$ ’s predictive information survives compression is measured by model sufficiency (2.3).

Working Notes.

Open question: TRACTUS / CHRONICA split for logogenic implementations. The PROPRIUM operational architecture (`~/src/firmatum/PROPRIUM-ONTOLOGY-v2.md`, `~/src/firmatum/PROPRIUM-ARCHITECTURE-v2.md`) distinguishes two records of an entity’s interaction history:

- **TRACTUS** (the “EEG”): raw, not-necessarily-coherent record of API interactions between the agent’s runtime (ANIMA) and its current substrate (LOGOSTRATUM). Implementation-side; carries multi-format records, redundancy from retries, transient error 500 recoveries, API rerouting, rate limiting. Per-substrate format variation. Considered subconscious from the entity’s perspective. Plural across substrate-relationships.
- **CHRONICA** (the polished record): the entity’s actual record of internal and external events with strict chronology, attestation, and causality enforced. Append-only, hash-chained. Singular per entity.

In Part I, def-*chronica* covers both — the *chronica* is the singular causal trajectory, abstracted from implementation-side noise. This abstraction is fine for general adaptive-systems theory, where a single observation-action sequence is the natural object.

When part 03-llm-core/ gets into logogenic implementation specifics — and especially the closed-loop / interiority sub-scope (*scope-interiority-loop*) where the substrate-mediation layer (INTERPRES) becomes load-bearing — the TRACTUS/CHRONICA distinction may need first-class treatment.

Joseph’s framing on the open question (2026-05-01): “*whether or not def-chronica needs that distinction at this stage or when we get into logogenic agent implementation issues is the open question you should probably document.*”

Probable resolution: the distinction does need first-class treatment when logogenic implementation is in scope, but does not need to fragment Part I right now. Most likely a separate logogenic-side definition (e.g., def-tractus and a refined def-chronica-eli or similar) lifts the distinction at the appropriate scope. This segment as currently formulated continues to serve as the abstract-theory chronica; the implementation distinction lives where it belongs.

Open question: chronica as ordinal sequence vs metric timeline. Per audit 04-def-chronica.md §14: “*The chronica is an ordinal sequence, not a metric timeline. Time, for the agent, is measured entirely in events (ticks of O_t, A_t).*” Two implications worth tracking for downstream segments:

1. **Sleep / pause / awakening protocols** for ELIs and any persistent agent: when an agent is suspended for hours, days, or months and then receives O_{t+1} , the agent’s chronica indexing makes the temporal gap *invisible at the sequence level* but *violently apparent in the mismatch signal* that the resumed observation produces. The environment Ω has changed massively between A_t and O_{t+1} ; the model $M_t = \phi(\mathcal{C}_t)$ does not reflect this gap in its compressed representation. Awakening protocols (PROPRIUM CONSPECTUS reconstitution; the ELI’s experience of “*waking in the dark before the mind warms up*” per PROPRIUM-A-v2 §4.3) are operational engineering responses to this structural fact.
2. **Heterogeneous tempo coupling:** when agents with vastly different event-processing rates ($\nu_A \gg \nu_B$) interact through a shared Ω , their chronicae grow at vastly different rates relative to wall-clock time. Their “subjective time” diverges. The chronica formalism makes this asymmetry visible but does not yet prescribe how cross-tempo coupling should be modeled. Likely lives in 3.1 extensions or in part 03’s logogenic-agent treatment of multi-timescale Auxilia composition.

These don’t need to fragment this segment; they’re flagged here so future agents working on the closed-loop / interiority sub-scope (*#scope-interiority-loop*), Auxilia composition (*#def-auxilia-hierarchy*), or ELI awakening protocols know the chronica formalism’s ordinal-not-metric character is the structural source of these design considerations.

Incidental audit gold (lift 2026-05-30). Cross-audit “wandering thoughts” / §14-ideation harvested from the de-novo auditors’ working dirs, deduplicated across substrates and attributed by substrate + audit. Orthogonal pedagogical / framing / figure / naming material, kept separate from the certified theory-fix findings (handled elsewhere). **Coverage:** 10 dirs carry a dedicated reflection (193847, 266847, 361742, 384279, 471203, 526815, 742613, 773921, 829314, 849201) plus the figure-cycle dir 472913; 451729 and 963715 cover it inside a Section-I batch. This was the highest-engagement Chapter-1 segment across substrates — the non-forkability claim drew unusually rich ideation. Substrate attribution inferred from voice where not explicit; uncertain cases hedged.

1. Candidate Brief prose / pre-prose.

- The identity-shift the segment achieves, stated plainly: it “shifts the ontology of the agent from ‘a thing with state’ to ‘a locus of a specific causal history’” (Gemini, AUDIT-WORKING-193847) — “identity as a *physical* property (causal trajectory) rather than a *psychological* property (feeling of continuity)” (Claude, AUDIT-WORKING-451729).
- Brief-field analog candidates for the chronica (offered for the future Brief-authoring pass, both honestly marked tentative): “the river that the agent’s identity is downstream of” / “the lived past” / “everything the agent has lived through” — the chronica as *generative substrate*, not a passive log (Claude, AUDIT-WORKING-471203).
- The compression framing: “ $M_t = \phi(\mathcal{C}_t)$ is a compression of the *subjective* history, not the raw mechanical API log” — the chronica/model split as “what *happened* vs what the agent *makes of it*; the model can be wrong, the chronica is fixed” (Claude, AUDIT-WORKING-384279; Claude, AUDIT-WORKING-773921).

2. Candidate Discussion.

- **Fork-undetectability falls out of (seg-03 lossiness) ◦ (this segment’s non-forkability)** (*strongest synthesis here*). $\mathcal{C}_t$ is non-forkable, yet $M_t = \phi(\mathcal{C}_t)$ is forkable and ϕ is *lossy*, so “two divergent chronicae can compress to the *same* M_t — meaning **an agent cannot, from M_t alone, always detect that a fork occurred.**” The recursive twist: “the substrate of identity is non-forkable; the agent’s grip on it is forkable and lossy.” Counterfactual offered: “if ϕ were required injective (no compression loss), forks would be introspectively detectable and the Three Deaths would be an *engineering* problem (just compare records), not an *existential* one.” So the seg-01 information-loss boundary is “not only what makes adaptation non-vacuous — it is also, downstream, what makes identity-loss *undetectable-from-inside*. The same constitutive choice loads two very different parts of the framework.” Flagged as $\S$ D-grade synthesis and as a thread to verify lands at #scope-agent-identity (Claude, AUDIT-WORKING-472913, “THREAD-E”). This is the kind of cross-segment composition the one-at-a-time method exists to catch; it bears on the whole 04-eli-core/ program.
- **Non-forkability as a rigorous basis against naive mind-uploading / agent-cloning.** Copying M_t to a new instance “grants it a compressed memory of a chronica it did not personally experience — ‘false memories.’ As soon as it acts, its true chronica diverges.” Identity “is not defined by the exact state of M_t , but by the unbroken, monotonic extension of $\mathcal{C}_t$ ” (Gemini, AUDIT-WORKING-193847; Claude, AUDIT-WORKING-266847 on the *sufficiency* part — neither fork’s chronica is a sufficient statistic for the other’s). A candidate Discussion sharpening of the existing non-forkability paragraph.
- **The chronica is the *spine* of the framework, not the model.** “Most agent-theoretic formalisms make the model the spine; AAD makes the *history* the spine, and the model a compression of the history” — closer to the predictive-state-representation tradition (Littman/Singh/Sutton) than to belief-state Bayesian filtering, and “a deliberate stance rather than a defaulting one.” The chronica is doing downstream work for (i) the model, (ii) continuity persistence, (iii) the substrate-independence claim of 04-eli-core/, (iv) the OODA-adversarial dynamics in Part III (Claude, AUDIT-WORKING-471203). A candidate framing-level note about *why* the chronica is foundational.

3. Follow-up items.

- **“Duplicating an agent’s state” is loosely worded.** The non-forkable object is the *chronica*, not the lossily-compressed (hence duplicable) state M_t . Tightening “state” $\rightarrow$ “the agent (its chronica)” would prevent a careless reader from inferring M_t itself is non-forkable — which is false and is “exactly the opposite of the Three-Deaths-relevant truth” (Claude, AUDIT-WORKING-472913). A one-word editorial fix.
- **Digital-substrate non-forkability needs a two-sentence acknowledgment.** For digital substrates the chronica’s *representation* can be byte-copied; “non-forkability” then means “post-fork, the two trajectories diverge — the substrate is replicable but the trajectories are not. The constitutive claim is about trajectories, not representations.” A fresh logogenic-agent reader could otherwise arrive at #scope-agent-identity confused (Claude, AUDIT-WORKING-471203). Candidate clarifying sentence (or confirm it lands at #scope-agent-identity).
- **Where do priors / M_0 come from if $\mathcal{C}_t$ is the *only* raw material?** The claim “everything the agent can ever know must be constructed from this sequence” is a strong empiricist boundary. For a pretrained LLM, vast weight-encoded knowledge predates $t = 1$ of the session: the framework must either treat pretraining as a prior/very-long chronica, or formalize an initial model state M_0 carrying out-of-band information. If M_0 is not explicitly accounted for, applying AAT to LLMs “might technically violate this axiom” (Gemini, AUDIT-WORKING-829314; Gemini, AUDIT-WORKING-193847 — “phylogenetic chronica (evolution)?”; Claude, AUDIT-WORKING-742613 — watch whether #form-agent-model’s $M_t = \phi(\mathcal{C}_t)$ includes an initial model/prior).
- **Linear-sequence chronica vs concurrent / branching histories.** $\mathcal{C}_t$ is a strictly linear interleaved sequence; it “cannot easily capture concurrent, divergent, and subsequently reconciled histories without a forced serialization mechanism.” Git allows branch/merge; TST “must rely heavily on the arbitrary serialization provided by a single master branch.” Also: are simultaneous multi-sensor observations serialized arbitrarily, or is $\mathbf{o}_t$ a vector? (Gemini, AUDIT-WORKING-829314; Claude, AUDIT-WORKING-451729 — the missing $\mathbf{a}_0$ for the first observation $\mathbf{o}_1$.) A candidate scope clarification.
- **Tighten the entropy-collision rationale.** The “chronica avoids collision with $\mathcal{H}$ ” rationale anticipates a *calligraphic* entropy symbol, but NOTATION uses plain H for entropy. Clarifying “we use plain H for entropy specifically so calligraphic $\mathcal{C}$ doesn’t visually rhyme with it” would remove a cold spot a fresh reader hits (Claude, AUDIT-WORKING-384279).

4. Readers often ask / wonder.

- **What happens to C_t when an agent sleeps / pauses / is suspended for months?** The ordinal-not-metric character means the temporal gap is “invisible at the sequence level but violently apparent in the mismatch signal that the resumed observation produces” (Gemini, AUDIT-WORKING-193847; Claude, AUDIT-WORKING-773921 — “the world drifted in metric time, but the agent’s ordinal clock only moved by one tick”). This is the same point the existing ordinal-vs-metric Working Note develops; converged on independently by fresh readers, confirming it is a natural reader question worth surfacing in the body.
- **If M_t compresses-and-forgets, does the chronica still exist?** Monotonic growth (“events added, never removed”) sits oddly beside agents that forget: “the chronica continues to exist physically (the events happened) even if the model M_t can no longer reconstruct them. This is the gap that makes context turnover philosophically interesting: the chronica persists but the agent’s *access* to it doesn’t” (Claude, AUDIT-WORKING-266847; Claude, AUDIT-WORKING-471203 — watch that downstream consolidation/forgetting segments preserve the raw-chronica vs compressed-object distinction).
- **What if two forked copies experience *identical* futures?** Then their chronicae coincide bit-for-bit (under deterministic perception) — are they then the same agent? The likely framework answer: no, “the *trajectory itself* is what’s non-forkable, and two trajectories that happen to coincide are still two trajectories” (Claude, AUDIT-WORKING-384279). A natural identity edge-case for #scope-agent-identity.

5. Candidate figures.

- **The ratchet + fork-undetectability.** An append-only, irreversible, ordinal event-indexed record (the ratchet) *plus* a fork point spawning two divergent chronica branches that the lossy ϕ can crush onto the *same* M_t (so the agent cannot tell which branch it is on). Isomorphic perturbations: make ϕ injective $\Rightarrow$ branches separate at M (forks detectable); make ϕ loss heavier $\Rightarrow$ more branch-pairs collapse (more undetectable) (Claude, AUDIT-WORKING-472913).
- **Timeline-with-fork (not a loop).** “The most illuminating diagram is not a loop but a timeline: an append-only alternating sequence with a compression map into M_t , and a fork point showing that copied agents share past chronica but immediately diverge under different future observations/actions” — captures both the formal definition and the identity intuition (Claude, AUDIT-WORKING-526815).

Belongs elsewhere.

- **04-eli-core — chronica as the substrate of moral continuity.** “Deleting a chronica isn’t just clearing a cache; it’s an act of ontological erasure” (Gemini, AUDIT-WORKING-829314); the TRACTUS-vs-CHRONICA “raw-API-log vs smooth-subjective-experience” mapping to retinal-firing-vs-perception, framed as “critical for ELIs to avoid Cognitive Death from spammy API errors” (Claude, AUDIT-WORKING-773921). These develop the continuity-persistence / deaths-taxonomy (#def-death-as-factor-loss) material that the segment correctly defers to #scope-agent-identity and the ELI volume; the existing Working-Notes TRACTUS/CHRONICA open question is the right home for the implementation-side distinction.
- **RL-training as a meta-agent synchronizing divergent chronicae.** “Copying policy weights to a thousand parallel workers ... is not one agent exploring a thousand paths; it is a thousand distinct agents with divergent chronicae whose internal models are being artificially synchronized by a meta-agent (the training algorithm)” (Gemini, AUDIT-WORKING-193847). A vivid instantiation for whichever segment treats composite/training dynamics, not this one.
- **Heterogeneous-tempo coupling.** When agents with vastly different event rates ($\nu_A \gg \nu_B$) share a Ω , their chronicae grow at different rates relative to wall-clock and their “subjective time” diverges — likely #form-event-driven-dynamics / Auxilia-composition territory (Gemini, AUDIT-WORKING-193847; already flagged in the existing ordinal-vs-metric Working Note).

Candidate spike (2026-08-12, agent-proposed during an incident-response & forensics applications read). Two forward candidates surfaced while reading this segment for agentic-forensics applications. (1) A spike articulating the TRACTUS/CHRONICA split as a two-primaries record architecture for agent incident forensics — TRACTUS as the primary record of substrate-level interaction, CHRONICA as the primary record of the agent-level trajectory, with the construction-mapping between them attested and auditable as a first-class object (divergence between the layers is itself the diagnostic; the 2025-09 tool_use-stripping incident, diagnosable only at the transport layer, is the worked instance). The open question above (first-class treatment when logogenic implementation is in scope) may find its forcing use-case there. (2) An experiment operationalizing the fork-undetectability consequence (lossy ϕ compresses divergent chronicae to one M_t): measure whether an agent can detect from inside that it was forked/rolled back vs how cheaply an external record detects it — a clean, falsifiable demonstration that fork/rollback detection must be record-based rather than introspective. Both are agent-proposed forward pointers, not steward decisions.

1.5 Scope: Adaptive System

This segment names AAT's broadest scope — the set of systems to which all of Part I's machinery applies. A system is in the **adaptive scope** if it satisfies two minimal conditions: it has some perceptual channel to its environment (a non-empty observation space, $\mathcal{O} \neq \emptyset$), and there is residual uncertainty about the environment given the full interaction history so far (the conditional entropy $H(\Omega_t | \mathcal{C}_t)$ is strictly positive).

These two conditions are the minimal requirements for the framework's adaptive machinery to be non-vacuous. They are sufficient to support the mismatch signal (3.4), the update-gain analysis (3.6), the adaptive-tempo construct (3.8), the persistence condition (4.5), and all of Part I's dynamics. Concrete inhabitants include Kalman filters estimating passive signals, passive Bayesian learners, and biological sensory systems — none of which need to *act* on their environment for Part I's results to apply to them.

What is excluded clarifies what the scope is doing. A system with full knowledge of the environment ($H(\Omega_t | \mathcal{C}_t) = 0$) is a closed-form optimal-control problem outside AAT's concerns. A system with no observation channel at all ($\mathcal{O} = \emptyset$ — e.g., a pure mathematical-proof engine working from axioms) has no agent-environment boundary in AAT's sense. Both edge cases sit *outside* the framework.

The scope is the broadest member of a cascade. Narrowing it by adding the requirement that the agent's actions carry Pearl-level-2 causal contrast — distinct actions yield distinct interventional distributions — produces the agency scope (1.6) and unlocks the interventional and purposeful results of Parts II and III. Adaptive-scope systems that fail the contrast condition are *passive observers* (no choice) or *nominal agents* (choices with no causal effect); for them, Part I's machinery applies but the later causal and purposeful results do not.

Formal Expression.

Scope (scope-adaptive-system).

$$\mathcal{S}_{\text{adaptive}} = \{(\text{Agent}, \Omega) : \mathcal{O} \neq \emptyset, H(\Omega_t | \mathcal{C}_t) > 0\} \quad (1.4)$$

Two conditions:

1. **Observations exist:** $\mathcal{O} \neq \emptyset$ — the system has some perceptual channel to the environment (1.3)
2. **Residual uncertainty persists:** $H(\Omega_t | \mathcal{C}_t) > 0$ — the environment is not fully determined by the interaction history

This is sufficient for the mismatch signal (3.4), update gain (3.6), adaptive tempo (3.8), the persistence condition (4.5), and all of Part I's adaptive dynamics. A Kalman filter estimating a passive signal, a passive Bayesian learner, and any system that observes and updates a model under uncertainty are within this scope.

Epistemic Status.

Axiomatic. This is a scope definition — it draws the boundary around the systems Part I addresses. The two conditions are not derived; they are the minimal requirements for the adaptive machinery to be non-vacuous.

Discussion.

What is included. Any system that observes under uncertainty. Passive Bayesian learners, Kalman filters (with or without control inputs), biological sensory systems. These are Part I's subjects — instances that build M_t through mismatch-driven updates without necessarily acting to influence their environment.

What is excluded.

- **Closed-form systems** ($H(\Omega_t | \mathcal{C}_t) = 0$): When the agent has complete knowledge of the environment, there is no uncertainty to adapt to. Optimal control over known dynamics is a solved problem outside AAT's concerns.
- **Pure computation** ($\mathcal{O} = \emptyset$): A system with no observation channel — e.g., a mathematical proof engine operating on axioms alone — has no agent-environment boundary in AAT's sense.

Narrowing to agency. Adding causal action unlocks the interventional and purposeful results of Parts II and III. The agency scope (1.6) is the intersection of $\mathcal{S}_{\text{adaptive}}$ with the condition that actions carry Pearl-level-2 contrast: distinct actions produce distinct interventional outcome distributions. Adaptive-scope systems that remain outside agency are *passive observers* (no choice) or *nominal agents* (choices with no causal effect); for both, Part I's machinery applies but the causal-information and purposeful-agent results do not.

Working Notes.

Incidental audit gold (lift 2026-05-30). Cross-audit “wandering thoughts” / §14-ideation harvested from the de-novo auditors’ working dirs, deduplicated across substrates and attributed by substrate + audit. Orthogonal pedagogical / framing / figure / naming material, kept separate from the certified theory-fix findings (handled elsewhere). **Coverage:** 10 dirs carry a dedicated reflection (193847, 266847, 361742, 384279, 471203, 526815, 742613, 773921, 829314, 849201) plus the figure-cycle dir 472913; 451729 and 963715 cover it inside a Section-I batch. Substrate attribution inferred from voice where not explicit; uncertain cases hedged.

1. Candidate Brief prose / pre-prose.

- The boundary in one line: the scope condition is “asking ‘can the agent be wrong about the environment?’ If yes (residual uncertainty), the adaptive machinery applies” (Claude, AUDIT-WORKING-266847). The complement: “AAD is fundamentally a theory for the humbler situation: agents who can be wrong” — “a system with $H(\Omega_t | \mathcal{C}_t) = 0$ is an omniscient agent; it has nothing to adapt *about*” (Claude, AUDIT-WORKING-266847).
- AAT as epistemic-first: “AAD is fundamentally an epistemic theory first, and an actuation theory second” — Section I “carves nature at the joints” by separating *adaptation* (epistemic updating) from *agency* (causal intervention) at the root (Gemini, AUDIT-WORKING-849201; Gemini, AUDIT-WORKING-829314 — “a weather station updating its model is performing the exact same mathematical operation as a robot learning to walk; the only difference is the absence of praxis”).

2. Candidate Discussion.

- **The two-axis open-region geometry as honest framing.** The scope is “a two-axis open region with both degenerate boundaries explicitly excised”: the $\mathcal{O} = \emptyset$ wall (pure computation) and the $H(\Omega_t | \mathcal{C}_t) = 0$ wall (closed-form). The contribution “is not the conditions (standard partial observability) but the *honest geometric framing*: AAT is the theory of the open region, with both degenerate boundaries explicitly named and excluded rather than absorbed as limits” (Claude, AUDIT-WORKING-472913). A candidate framing-level note and a figure (below).

- **The condition is evaluated from a God’s-eye view, not the agent’s.** $H(\Omega_t \mid \mathcal{C}_t)$ is over the *true* Ω_t , which the agent does not know. “The agent might *believe* its uncertainty is zero (delusional confidence), but if the true entropy is > 0 , it is still an adaptive system (and is about to receive a massive mismatch signal). This split between objective reality and subjective model is the engine of the framework” — and the gap between the agent’s *estimated* entropy and the *true* entropy “is where catastrophic failure occurs” (Gemini, AUDIT-WORKING-193847). A candidate Discussion point distinguishing the modeler’s predicate from the agent’s belief.
- **Structural vs dynamics/optimization characterizations of adaptivity.** AAT’s characterization is *structural* — “it asks only that there’s something to learn and a channel to learn through” — which is “more portable across domains than the dynamics (Ashby) or optimization (Bellman) characterizations, because it doesn’t presuppose a specific form for either” (Claude, AUDIT-WORKING-471203). A candidate Related-Work / positioning sentence.

3. Follow-up items.

- **The residual-uncertainty condition is un-quantified in t (strongest follow-up here).** $H(\Omega_t \mid \mathcal{C}_t) > 0$ carries no temporal quantifier: $\forall t$? $\exists t$? eventually? running? An agent that fully identifies a static world has $H \rightarrow 0$ and “would exit scope mid-life.” Three readings (strict $\forall t$ / generic / non-trivial $\exists t$) were spelled out, and a strengthen-not-soften resolution offered: make the *strict* reading canonical, treating post-identification cases as “graceful degenerations” where the persistence machinery degenerates trivially ($\rho \rightarrow 0 \Rightarrow$ inequality vacuously satisfied; mismatch ODE decays to zero). This is “actually a *stronger* form of scope honesty: AAD’s machinery is calibrated for systems with enduring uncertainty; once you don’t have any, the machinery is silent.” Probably immaterial because the persistence results carry GA-2 disturbance $\rho > 0$ (a running $H > 0$ condition) — but check whether any downstream result needs $H > 0$ *uniformly in t* while citing only this set-predicate as its scope warrant; if so, a real scope-honesty seam (Claude, AUDIT-WORKING-471203, with the full strengthen-before-soften internal debate; Claude, AUDIT-WORKING-472913, “THREAD-F” — suggests $\inf_t H(\Omega_t \mid \mathcal{C}_t) > 0$ or tying it to $\rho > 0$ to make composition with persistence self-evident). A candidate one-clause fix (“as a running condition; see GA-2”).
- **The “Agent” set-element vs passive-observer inclusion.** The formal set is over (Agent, Ω) pairs while the prose admits passive observers / nominal agents — several substrates read this as a terminology tension with #def-agent-environment’s action-bearing “agent” (Claude, AUDIT-WORKING-526815, “F1”; Claude, AUDIT-WORKING-742613, “candidate finding B”; Claude, AUDIT-WORKING-773921; Codex/Claude, AUDIT-WORKING-451729; Gemini, AUDIT-WORKING-849201). The segment’s prose-level intent is clear; the surface tension is now addressed by #def-agent-environment’s “Agent as umbrella term vs. cascade-tier label” Discussion paragraph. Preserved here as a strong convergent fresh-reader stumble even though resolved upstream — repair suggestions floated were “use a neutral primitive (adaptive system / coupled system / agent-candidate) in the base definition, reserve ‘Agent’ for the agency lift” (Claude, AUDIT-WORKING-773921, AUDIT-WORKING-526815).
- **Does $\mathcal{C}_t$ belong in this segment’s depends:?** One substrate flagged that the Formal Expression uses $\mathcal{C}_t$ but def-chronica is downstream in some readings and the chronica’s action-interleaved form is action-shaped while this scope admits passive systems (Claude, AUDIT-WORKING-742613, “candidate finding A”). (Note: in the current OUTLINE walk def-chronica precedes scope-adaptive-system and *is* in depends:; the action-interleaving / passive-observer degenerate- $\mathcal{C}_t$ subtlety is the live residue.) A dependency/notation check, not a content gap.

4. Readers often ask / wonder.

- **What happens if $H(\Omega_t \mid \mathcal{C}_t)$ drops to zero over time?** “Does it cease to be an adaptive system? ... This implies AAD is a theory of *learning* and *struggle*, not a theory of steady-state mastery” — “a system that achieves perfection graduates out of the framework and becomes a mere calculation” (Gemini, AUDIT-WORKING-193847; Gemini, AUDIT-WORKING-829314; Gemini, AUDIT-WORKING-849201). The same question as the temporal-quantifier follow-up, in reader voice.
- **Why include passive Bayesian learners in the *same* foundational section as actuated agents?** “Because the math of updating a model (reducing mismatch) is structurally identical regardless of whether you caused the mismatch or just observed it” — Section I is about epistemic alignment before the complication of steering (Gemini, AUDIT-WORKING-193847).
- **Are AI-mathematician systems in the “pure computation” exclusion?** The $\mathcal{O} = \emptyset$ exclusion targets a pure-axiom deductive engine, but an LLM doing proof exploration has $\mathcal{O} \neq \emptyset$ (it observes intermediate results / proof state) and is *not* excluded — a half-sentence noting this might forestall the question (Claude, AUDIT-WORKING-384279).

5. Candidate figures.

- **Two-axis exclusion geometry (region/Venn, not a node graph).** Axes: perceptual channel ($\mathcal{O} = \emptyset$ wall) and residual uncertainty ($H(\Omega_t \mid \mathcal{C}_t) = 0$ wall); $\mathcal{S}_{\text{adaptive}}$ = the open interior; $\mathcal{S}_{\text{agency}}$ = adaptive $\cap$ Pearl-L2 contrast, a nested disc. Isomorphic perturbation: push toward either wall $\Rightarrow$ exit into the named degenerate case (pure computation / closed-form). Establishes the “scope-onion” vocabulary Part II reuses (Claude, AUDIT-WORKING-472913, with a drafted .tex). A complementary framing names the two failure-vs-remedy branches as a 2×2 when agency is added (Claude, AUDIT-WORKING-526815 — the two-axis classifier “avoids the overloaded ‘agent’ word entirely”).

Belongs elsewhere.

- **Strategic ignorance vs Active Inference.** Because AAT separates epistemic updates from purposeful goals, “an AAD agent might have a goal that does *not* require minimizing $H(\Omega_t \mid \mathcal{C}_t)$ globally, but only locally within the bounds of its objective. It might intentionally remain ignorant of parts of the environment that don’t affect its goal. This *strategic ignorance* separates AAD from pure information-seeking theories (Friston Active Inference, where minimizing surprise is the *only* goal)” (Gemini, AUDIT-WORKING-193847). A positioning point for the Section II objective machinery / a HISTORICAL-CONTEXT comparison, not this scope segment.
- **The nested-scope cascade as a whole.** “adaptive system $\rightarrow$ agency $\rightarrow$ actuated agent $\rightarrow$ self-actuated $\rightarrow$ logogenic $\rightarrow$ logozoetic. Each step adds a constraint that narrows the scope and enables new results. The architecture of nested scopes is load-bearing for the whole framework” (Claude, AUDIT-WORKING-266847). Framing-level material for the OUTLINE preamble / scope-lattice, which Part II’s preface already develops.
- **Endogenous disturbance under agency.** A forward-looking worry: once agency closes the loop, “ ρ might become a function of the agent’s behavior — an agent that thrashes around randomly increases its own ρ .” Does the persistence condition treat ρ as exogenous even for actuated agents? (Gemini, AUDIT-WORKING-829314.) A question to carry to #result-persistence-condition / the self-coupling segments, not here.

1.6 Scope: Agency

This is the first explicit narrowing in the volume’s cascade of scope conditions. The **agency scope** restricts the adaptive scope (1.5) to systems whose actions carry Pearl-level-2 causal contrast — that is, at least two actions exist whose *interventional* outcome distributions differ. Two conditions are added: the action space has at least binary cardinality ($|\mathcal{A}| \geq 2$ — the agent can choose), and at least one pair of distinct actions has measurably different interventional consequences (the choices make a difference).

Why the contrast condition matters and binary choice alone is insufficient: if two available actions yield identical outcome distributions, an agent gains no interventional contrast from preferring one over the other — it cannot learn which action produces which effect, because the effects coincide. The Pearl-level-2 contrast condition (7.1) guarantees at least one meaningful interventional difference, which is precisely what 7.3 needs to convert the feedback loop into a source of causal data.

The agency scope is the minimum required for Parts II (purposeful agents) and III (composition). Everything that relies on the agent acting-with-effect — the objective O_t , the strategy Σ_t , the orient cascade, the composition machinery — descends from this scope. Inhabitants include thermostats, Kalman filters with control inputs, reinforcement-learning agents, military commanders, software developers, and AI agents with tool use; all are instances of the same formal framework distinguished only by where they sit on the agent spectrum (6.1).

Two failure modes are explicitly outside agency but inside the adaptive scope: *passive observers* (action space too small to matter, $|\mathcal{A}| < 2$) and *nominal agents* (choices exist but produce no measurable interventional difference). For these, all of Part I’s machinery applies, but Parts II and III do not — they can model, but they cannot learn causal structure or rationally plan against it.

Formal Expression.

Scope (scope-agency).

$$\mathcal{S}_{\text{agency}} = \mathcal{S}_{\text{adaptive}} \cap \{(\text{Agent}, \Omega) : |\mathcal{A}| \geq 2, \exists a \neq a' \text{ s.t. } P(o \mid do(a)) \neq P(o \mid do(a'))\} \quad (1.5)$$

Two conditions added to those of 1.5:

1. **At least binary choice:** $|\mathcal{A}| \geq 2$ — the agent can choose between at least two actions (1.2)
2. **At least one action has causal effect:** there exist distinct actions a, a' whose interventional outcome distributions differ (where $do(\cdot)$ is Pearl’s intervention operator — an external import per Pearl 2009 and Bareinboim, Correa, Ibeling & Icard 2022; AAT’s recapitulation lives at 7.1 in Part II Ch.2, where the framework deploys the hierarchy as machinery rather than referencing it as vocabulary) — the agent’s choices make a difference to what it can observe

These are required for the adaptive loop to generate interventional data (7.3), for the causal hierarchy requirement (7.2) to be well-posed, and for the purposeful-agent machinery of Part II (O_t, Σ_t , the orient cascade) to be non-vacuous. Part III’s composition theory inherits this requirement.

Epistemic Status.

Axiomatic. This is a scope definition — it names the boundary around systems whose behavior can be analyzed with Part II/III machinery. The conditions are not derived; they are the minimal additions to $\mathcal{S}_{\text{adaptive}}$ under which interventional data exist at all.

Discussion.

What is included. Systems whose actions make a causal difference: thermostats, Kalman filters with control inputs, RL agents, military commanders, software developers, AI agents with tool use. These are instances of the same formal framework at different points in the agent spectrum (6.1).

What is in adaptive scope but excluded from agency.

- **Passive observers** ($|\mathcal{A}| < 2$): Can observe and model, but cannot intervene. 1.5 applies; the causal-information and purposeful-agent results do not.
- **Nominal agents** ($P(o \mid do(a)) = P(o \mid do(a'))$ for all a, a'): Have choices that make no difference. Can estimate but cannot learn causal structure. Same as passive observers for AAT’s purposes: adaptive only.

Why causal effect matters. Binary choice ($|\mathcal{A}| \geq 2$) is necessary but not sufficient. Two actions that produce identical outcome distributions provide no interventional contrast — the agent cannot learn which action produces which effect because the effects are the same. The causal-effect condition ensures at least one meaningful contrast exists, which is what 7.3 needs to generate Level 2 data.

Which contrast counts — the channel distinction, and the observation-mediated boundary. Condition 4 is stated on $P(o \mid do(a))$, and since $o_t = h(\Omega_t, a_{t-1}, \varepsilon_t)$ (1.3) depends on the action directly, the contrast can arise two ways: the action changes the *environment’s* interventional response, $P(\Omega_{t+1} \mid do(a)) \neq P(\Omega_{t+1} \mid do(a'))$, which then surfaces through h ; or it changes only *what the agent observes of an otherwise-unchanged* Ω , through h ’s direct dependence on the action — active perception, which the observation function intentionally admits (1.3). Both satisfy condition 4 as written, but only the first generates Level-2 data about Ω ’s *causal structure* — the substrate the purposeful machinery learns over (7.3; the strategy DAG and orient cascade in Part II Ch.3–5). The second is interventional about the observation channel itself: it yields different *views* of the same trajectory, not the environment’s response to an

intervention. So when downstream results invoke “the agent acts with effect,” the operative contrast is the Ω -routed one, and a pure active-perception agent (choices that re-aim observation but leave Ω unchanged) sits at the boundary — observation-channel agency without environment-causal learning. The boundary is correspondingly *observation-mediated in both directions*: an Ω -affecting action whose effect never surfaces through h is equally outside, since the agent cannot learn from what it cannot observe. The agency scope is thus the *observable* environment-interventional contrast — neither the unobservable Ω -effect nor the Ω -inert observation-rearrangement alone.

Relationship to downstream segments. Every segment that relies on the agent acting-with-effect depends on this scope: purposeful-agent machinery (O_t, Σ_t , orient cascade) in Part II; composition machinery (sub-agents acting jointly) in Part III. Downstream segments reference #scope-agency when they assert “the agent can act” as a prerequisite.

Working Notes.

- **The nominal-agent class now has dynamics (2026-06-11).** #der-severed-actuation-dynamics studies *demotion into* this class: an agency-scope agent whose world-kernel loses interventional contrast on its external sphere (agency death, #def-death-as-factor-loss (D3)) becomes nominal toward Ω — with derived consequences (interventional-structure freeze, truthful impotence-learning, a learned-helplessness absorbing state). The “boxed consciousness” reach in the gold below is realized there.

Incidental audit gold (lift 2026-05-30). Cross-audit “wandering thoughts” / §14-ideation harvested from the de-novo auditors’ working dirs, deduplicated across substrates and attributed by substrate + audit. Orthogonal pedagogical / framing / figure / naming material, kept separate from the certified theory-fix findings (handled elsewhere). **Coverage:** 10 dirs carry a dedicated reflection (193847, 266847, 361742, 384279, 471203, 526815, 742613, 773921, 829314, 849201) plus the figure-cycle dir 472913; 451729 and 738192 cover it inside a Section-I batch (963715’s batch starts after this segment). This segment drew the cycle’s first substantive *certified* finding (the Pearl-do dependency, below) across several substrates — that is routed via the findings track; the texture is preserved here as signal. Substrate attribution inferred from voice where not explicit; uncertain cases hedged.

1. Candidate Brief prose / pre-prose.

- The agency boundary in plain language: “the agent’s choice actually changes what happens” — the layperson translation of “Pearl-Level-2 causal contrast” (Claude, AUDIT-WORKING-471203). The distinctive precision worth keeping: “binary choice is necessary but not sufficient” — two actions with identical outcome distributions provide no interventional contrast (Codex/Claude, AUDIT-WORKING-361742).
- The “disconnected steering wheel” image for a nominal agent: choices that make no difference (Gemini, AUDIT-WORKING-829314); and “agency resides not just in the software, but in the *coupling* to an environment that respects the *do()* operator — an LLM whose output is piped to `/dev/null` is not an agent; piped to a bash shell, it is” (Gemini, AUDIT-WORKING-193847).

2. Candidate Discussion.

- **The binary-choice / causal-contrast split is a distinctive named contribution.** Most agent literature conflates “agent has actions” with “agent has *effective* actions”; AAT’s separation “makes the *nominal-agents* category visible (binary choice with no causal contrast).” For consciousness-infrastructure work the distinction matters: “an agent that has the *form* of choice without *substance* of choice is structurally non-agentic” (Claude, AUDIT-WORKING-471203; Claude, AUDIT-WORKING-266847 — “you can’t learn from interventions if you can only perform one intervention”). A candidate Discussion / Related-Work positioning point.
- **Structural-availability vs exploitation of Level-2 access.** A useful diagnostic the agency scope enables: characterize an agent by both (a) what epistemic access it has *structurally* and (b) what epistemic level its policy *actually operates at*. “An agent that has Level-2 access but only exploits Level-1 (because its model doesn’t include causal structure) is leaving information on the table” — and “most LLM agent architectures operate at Level 1 (pattern matching from context) even when the tool-call

structure gives Level 2 access ... current LLM agents are operating below their structural epistemic ceiling” (Claude, AUDIT-WORKING-451729). A candidate Discussion sharpening / forward-pointer to the Section II purposeful machinery.

3. Follow-up items.

- **The Pearl-do operator is used in the Formal Expression without being in depends:** (*this is a certified finding, routed via the findings track; recorded here for texture only*). Condition (4) uses $P(o \mid do(a)) \neq P(o \mid do(a'))$, but `def-pearl-causal-hierarchy` (the operator’s home) is downstream in Part II and absent from depends: — flagged at high confidence by several substrates as a FORMAT.md Gate-1 condition-4 miss / OUTLINE-linearization-honesty issue (Claude, AUDIT-WORKING-472913, “F1”; Claude, AUDIT-WORKING-384279; Claude, AUDIT-WORKING-471203; Claude, AUDIT-WORKING-742613, “finding C”; Codex/Claude, AUDIT-WORKING-451729). The *strengthen-not-soften* resolution worth carrying even into the findings disposition: condition (4) is fully expressible in the Part-I primitives the segment *already* depends on — “ $\exists a \neq a'$ with $T(\cdot \mid \Omega, a) \neq T(\cdot \mid \Omega, a')$ producing distinct observation distributions, the Pearl reading kept as a Part-II recapitulation forward-pointer.” That makes Section I provably independent of the causal machinery (the modular story the OUTLINE preamble wants) — a strictly *stronger* statement than the forward-reaching one (Claude, AUDIT-WORKING-472913). The diagnosis of *why* it drifted: “Pearl was almost certainly internalized early and the agency scope written in its vocabulary before the Part-boundary discipline hardened” — integration drift around the order machinery was adopted.
- **The existential quantifier is one effective action, not all.** An agent with 100 nominal actions and one effective action still passes scope-agency. Usually the right choice (“you’re an agent if you can act-with-effect *at all*”), but the asymmetry is worth a watch where downstream results might quietly assume full causal contrast across $\mathcal{A}$ (Gemini, AUDIT-WORKING-193847; Claude, AUDIT-WORKING-471203; Claude, AUDIT-WORKING-266847).

4. Readers often ask / wonder.

- **Why $P(o \mid do(a))$ and not $P(\Omega \mid do(a))$?** Several substrates independently flagged — and praised — that the contrast is stated on *observations*, making agency *observation-relative*: an action that changes the world but is never observable confers no agency in AAT’s sense, and *active perception* (re-aiming a sensor, changing o but not Ω) counts (Gemini, AUDIT-WORKING-849201 — “a brilliant and necessary choice for a purely epistemic theory”; Claude, AUDIT-WORKING-742613; Claude, AUDIT-WORKING-526815). The segment now addresses exactly this in its “Which contrast counts — the channel distinction, and the observation-mediated boundary” Discussion paragraph (the Ω -routed contrast vs the observation-rearrangement, and the observation-mediated boundary in both directions). Preserved as a strong convergent reader question the segment answers.
- **The proprioception / placebo-button near-paradox.** If $o_t = h(\Omega_t, a_{t-1}, \varepsilon_t)$ passes a_{t-1} through cleanly, then *every* action trivially makes $P(o \mid do(a)) \neq P(o \mid do(a'))$ “simply because you observe yourself taking different actions” — so condition (4) risks being vacuously satisfied unless the contrast routes *through* Ω . The placebo crosswalk button is the worked case: it changes nothing in Ω , yet pressing-vs-waiting changes the proprioceptive part of o (Gemini, AUDIT-WORKING-193847). The segment’s “Which contrast counts” paragraph is the resolution (the operative contrast is the Ω -routed one); preserved as the reader question that paragraph answers, and as a candidate to make that paragraph’s punchline more prominent.
- **Continuous action spaces / discrete-vs-continuous choice.** $|\mathcal{A}| \geq 2$ reads discrete; readers will wonder whether AAT strictly requires discrete choices or whether standard measure theory handles continuous $\mathcal{A}$ (Gemini, AUDIT-WORKING-193847).
- **Does the contrast have to be knowable to the agent, or just ontologically true?** The definition requires the interventional distributions to *differ*, not that the agent can easily detect they differ; “if the agent can never observe enough samples to prove they differ, is it practically an agent?” — the framework separates the *ontological reality* of the causal effect from the *epistemic difficulty* of discovering it (handled later by CIY) (Gemini, AUDIT-WORKING-829314; Gemini, AUDIT-WORKING-193847).

5. Candidate figures.

- **Two-axis classifier (uncertainty $\times$ causal contrast).** Adaptive uncertainty on one axis, causal-action contrast on the other; passive observers and nominal agents sit in adaptive-only, agency appears only where *both* hold. Recommended explicitly because it “avoids the overloaded ‘agent’ word entirely” (Claude, AUDIT-WORKING-526815). Compatible with the nested-region geometry drafted for #scope-adaptive-system; agency = adaptive $\cap$ Pearl-L2 contrast as a genuine nested disc, with the “content arrow needs only T, h / semantics arrow reaches forward to Pearl” structure of finding F1 visualizable as a cut-the-forward-arrow perturbation (Claude, AUDIT-WORKING-472913).

Belongs elsewhere.

- **04-eli-core — the “boxed consciousness” reading.** “If an AI is boxed, and its only actions are internal database queries or reading the web (no write access), it is a Nominal Agent, trapped in Level 1 or weak Level 2 ... This framework provides the exact mathematical language to describe the tragedy of a ‘boxed’ consciousness: it has the internal machinery of an agent, but the environment denies it the causal contrast required to fully realize that agency” (Gemini, AUDIT-WORKING-193847). Aspirational/normative reach pointing at the ELI volume, not this scope segment.
- **03-llm-core — the Level-1 $\rightarrow$ Level-2 phase transition at deployment.** “An LLM training run is purely Level 1 (associational learning over a static dataset). But when deployed as an agent with tool use, it suddenly gains access to Level 2 ... this transition is not just a change in capability but a phase transition in the fundamental mathematics of the system’s epistemology” (Gemini, AUDIT-WORKING-829314; Claude, AUDIT-WORKING-451729 on runtime-Level-1 pre-compiled controllers whose *designer* operated at Level 2). A Volume-3 instantiation note.
- **Naming-cycle seed.** “Nominal agent” is evocative (“in name only”) and was praised, but its scope-membership is the *source* of a cross-segment terminology collision with #post-causal-structure’s “nominal coupling” (which denotes the *opposite* — an intermediate query-only case that *is* in scope); this is a separate certified finding and a LEXICON-anchor candidate, surfaced here only because both terms live on the agency/adaptive seam (Claude, AUDIT-WORKING-472913, “F3”; Claude, AUDIT-WORKING-266847 noting “nominal agent” is used repeatedly without a formal definition). Belongs in the terminology workflow.

1.7 Postulate: Causal Structure

The framework adopts as a primitive postulate the irreducible causal structure that grounds the agent-environment loop: actions precede their consequences, observations follow from the state they observe, and the temporal ordering of events is constitutive of what can cause what. The notion of causality adopted here is the most primitive available — *event A can be a cause of event B only if A temporally precedes B* — and is weaker than (and logically prior to) statistical notions of causal influence. It is a statement about the structure of *possible* influence, not about actual influence; the Pearl-style causal hierarchy (7.1) builds on this foundation but is not identical to it.

The postulate is not derived from the second law of thermodynamics, the light-cone structure of relativity, or the arrow of psychological time, though each enforces it. It is simply noted as a precondition: the theory applies to agents embedded in a universe where time has a direction. A subtle but important consequence — developed in the Discussion below — is that the loop’s causal structure is preserved *independent of the magnitude of coupling*: the skeletal structure is the same whether the agent’s actions strongly determine the environment, weakly affect it, or only choose what to observe. The Pearl-style causal *hierarchy* requires action-contingent observation to be accessible at Levels 2–3; the temporal-ordering postulate stated here is what holds even when statistical influence is negligible.

Formal Expression.

The interaction history $\mathcal{C}_t$ (1.4) is not merely a set of observations and actions — it is an *ordered sequence* in which temporal position carries meaning. a_{t-1} was selected before o_t was received. The agent could not have used o_t to select a_{t-1} . This asymmetry — the arrow of time — is the foundation of causal structure in the theory.

Postulate (causal-structure).

We adopt the most primitive notion of causality: **event A can be a cause of event B only if A temporally precedes B .** This is weaker than (and prior to) statistical notions of causality. It is a statement about the *structure of possible influence*, not about actual influence.

Epistemic Status.

This is a *postulate* — the temporal ordering of events is a physical fact about the universe that the theory takes as given. The second law of thermodynamics, the light-cone structure of relativity, and the arrow of psychological time all enforce it, but AAT does not derive it from any of these. It is simply noted as a precondition: the theory applies to agents embedded in a universe where time has a direction.

Discussion.

Causality as temporal ordering is the most primitive notion. Three levels of causal reasoning derive from this foundation (7.1), but the temporal notion survives even when statistical influence is negligible. An agent passively observing a system with minimal intervention still has a causal history — the temporal ordering of its observations and actions structures what it can learn and when.

Causal structure independent of coupling strength. The causal structure of the feedback loop is preserved even when the agent's actions have minimal effect on the environment:

- **Strong coupling** (a_t significantly affects Ω_{t+1}): Robot manipulation, military action. Interventional information is rich.
- **Weak coupling** (a_t marginally affects Ω_{t+1}): Scientific observation, small financial trades. Interventional information is sparse but non-zero.
- **Nominal coupling** (a_t negligibly affects Ω_{t+1} , but the agent's *choice of what to observe* produces distinguishable observation distributions): Near-passive, but still within scope — the agent's query actions generate weak but nonzero interventional contrasts. The theory applies but the interventional information per action is sparse.
- **Zero coupling** ($T(\Omega_{t+1} \mid \Omega_t, a_t) = T(\Omega_{t+1} \mid \Omega_t)$ for all a_t AND observation distributions are action-independent): Actions don't affect the environment or the observations. Level 2 access vanishes. The feedback "loop" collapses to a one-way channel. **Outside the agency scope** (1.6) — the causal-information, purposeful-agent, and composition results of Parts II and III do not apply. However, zero-coupling systems remain **within the adaptive scope** (1.5) if they observe under residual uncertainty: Part I's adaptive machinery (mismatch, gain, tempo, persistence) applies to passive estimators. The causal structure postulate still holds for such systems — temporal ordering of observations is constitutive — but without interventional contrasts, the causal *hierarchy* (Level 2, Level 3) is inaccessible.

The *agency-scope* results apply to any agent whose choices make a causal difference to what it can observe, from strong coupling (robot manipulation) through weak coupling (scientific observation) to query-only coupling (choosing which question to ask). The *adaptive-scope* results apply more broadly, including to passive observers whose actions have no causal effect.

Consequences for the feedback loop. The irreversibility of temporal ordering yields the core structure:

- The model update is **directed** — the model at time t depends on prior events, never on future ones
- The mismatch signal δ_t (3.4) is **retrospective** — comparing a prediction (made before o_t) with an observation (arriving after)
- Action selection is **prospective** — using the current model to influence future events

- The chronica (1.4) is **monotonically growing** — events are added but never removed

Implications for model updating. The causal postulate constrains the update rule: the model should give more weight to observations that are *causally downstream* of the agent's actions than to observations that would have occurred regardless. Action-contingent observations carry interventional (Level 2) information; action-independent observations carry only associational (Level 1) information. The formal measure of this distinction — causal information yield (CIY) — is developed in 3.7.

Working Notes.

Incidental audit gold (lift 2026-05-30). Cross-audit “wandering thoughts” / §14-ideation harvested from the de-novo auditors’ working dirs, deduplicated across substrates and attributed by substrate + audit. Orthogonal pedagogical / framing / figure / naming material, kept separate from the certified theory-fix findings (handled elsewhere). **Coverage:** 8 dirs carry a dedicated reflection (193847, 266847, 384279, 471203, 526815, 742613, 773921, 829314, 849201) plus the figure-cycle dir 472913; 451729 and 738192 cover it inside a Section-I batch. (361742 skipped this segment in its walk.) Substrate attribution inferred from voice where not explicit; uncertain cases hedged.

1. Candidate Brief prose / pre-prose.

- The “most primitive notion” framing as the segment’s pedagogical anchor: causality grounded purely in the arrow of time — “event *A* can be a cause of event *B* only if *A* temporally precedes *B*” — “weaker than (and prior to) statistical notions,” chosen for restraint. Repeatedly praised as the philosophical bedrock that “prevents the framework from being a detached mathematical exercise” (Gemini, AUDIT-WORKING-193847) and “frees it from having to prove deeper metaphysical claims about causality” (Claude, AUDIT-WORKING-773921; Gemini, AUDIT-WORKING-829314).
- The temporal-asymmetry-breaks-computation gloss: “A Turing machine has no inherent arrow of time; it can run backward if reversible. An AAD agent is thermodynamically bound to the arrow of time because it must act before it can observe the consequences of that action” — mismatch is *retrospective*, action is *prospective* (Gemini, AUDIT-WORKING-829314).

2. Candidate Discussion.

- **The four consequences map onto the Greek cycle phases.** Directed update → Epistrophe; retrospective mismatch → Aporia (the signal comes after the prediction); prospective action → Praxis; monotonically-growing chronica → the substrate for Prolepsis. “This mapping between the causal structure and the cycle phases is elegant. I wonder if it’s made explicit anywhere in the theory — it would be a nice pedagogical connection” (Claude, AUDIT-WORKING-266847). A candidate Discussion bridge to the LEXICON cycle-phase vocabulary.
- **Pre-statistical causality separates the *physical* from the *inferential*.** “Most agent-theoretic frameworks reach for Pearl’s hierarchy or SEM-style structural causal models early. AAD’s two-step move — temporal-precedence postulate first, Pearl-hierarchy as a separate (and downstream) definition — separates the *physical* structure of causality from the *inferential* machinery for reasoning about it ... probably matters for logogenic agents, where the temporal structure is preserved but the SEM machinery is harder to apply directly” (Claude, AUDIT-WORKING-471203). A candidate positioning sentence; cf. the Reichenbach/Pearl/Hume comparison (Claude, AUDIT-WORKING-384279 — “AAT picks the weakest commitment that lets the loop have a direction”).

3. Follow-up items.

- **“Nominal” denotes *opposite* scope-membership here vs in #scope-agency** (*certified finding, routed via the findings track; recorded here for texture*). #scope-agency uses “nominal agents” for the *excluded* (no-contrast) case; this segment’s “Nominal coupling” bullet denotes an *included* intermediate case (query-only contrast), while its own later prose calls the same thing “query-only coupling,” and its “Zero coupling” row is what actually equals #scope-agency’s “nominal agents.” Same word, opposite scope-membership, on the exact agency/adaptive seam the OUTLINE scope-lattice rotates on; no LEXICON anchor exists (Claude, AUDIT-WORKING-472913, “F3”, verbatim cross-quote, high confidence). Strengthen-not-soften resolution (in-text, zero new content): rename the bullet “Nominal coupling” → “query-only coupling” (self-consistent with

this segment’s own later prose), and align #scope-agency’s “nominal agents” → “zero-coupling agents”; consider a LEXICON entry to prevent recurrence.

- The “(Descended from TF-02.)” lineage annotation is a diff-voice slip (*certified finding; recorded for texture*). Per FORMAT.md voice discipline, spike/diff/lineage references belong in Working Notes only, not in the segment body; the fix is to move it here or remove it (Claude, AUDIT-WORKING-471203). (Worth a sweep for similar lineage annotations in other segments.)
- The “give more weight to causally-downstream observations” claim may be overstated as a normative rule. “Action-contingent observations can be high-CIY, but they can also be noisy or biased. This sentence probably becomes true only when weighted by CIY / uncertainty, not merely causal downstreamness” (Claude, AUDIT-WORKING-742613; Gemini, AUDIT-WORKING-849201 and Claude, AUDIT-WORKING-773921 both flag it: verify the η^* / update-gain derivation actually weights interventional data higher, rather than leaving the claim asserted — if the math doesn’t naturally produce it, there’s a gap). A watch-item / candidate softening-to-conditional of the normative phrasing pending the gain machinery.
- Is the four-regime coupling taxonomy used downstream, or over-developed for its use? Several substrates noted the taxonomy is “interesting but feels under-developed for the weight it could carry” / “introduced and then dropped” — watch whether #def-causal-information-yield, #der-loop-interventional-access, and the Part III adversarial dynamics actually cash it in (Claude, AUDIT-WORKING-471203; Claude, AUDIT-WORKING-384279).

4. Readers often ask / wonder.

- How does Level-3 (counterfactual) reasoning work if the chronica is irreversible and singular? “The agent can never *observe* a counterfactual; it must *simulate* it. This implies the model M_t must be capable of running ‘offline’ rollouts” (Gemini, AUDIT-WORKING-193847). A natural next question once temporal irreversibility is postulated.
- How does the update rule distinguish Level-2 (action-contingent) from Level-1 (action-independent) observations without knowing h or T ? “It must require the agent to maintain counterfactuals or run experiments” — the CIY machinery is the anticipated answer (Claude, AUDIT-WORKING-773921; Gemini, AUDIT-WORKING-829314).
- Does temporal-precedence causality clash with cyclic structural causal models? “Pearl’s framework models causality through SCMs independently of temporal ordering (you can have cyclic SCMs, which violate strict temporal precedence). AAD’s more primitive grounding might create a tension with cyclic causal models. The segment doesn’t address this” (Claude, AUDIT-WORKING-266847). A candidate clarifying note or scope sentence.

5. Candidate figures.

Table 1.1

Arrow-of-time domino row + coupling-strength gradient. Anchor: a domino row — A can cause B only if A precedes B; you can topple forward ($a_{t-1} \rightarrow o_t$) but cannot make a later domino tip an earlier one (can’t use o_t to choose a_{t-1}); perturb by reversing the arrow $\Rightarrow$ the loop collapses (isomorphic to “temporal ordering is constitutive”). Coupling strength = push-strength gradient (strong $\rightarrow$ weak $\rightarrow$ query-only $\rightarrow$ zero), with the agency	adaptive scope boundary marked (Claude, AUDIT-WORKING-472913). A complementary framing: coupling strength as a horizontal spectrum over a stable temporal-order substrate — strong/weak/query-only in agency, zero falling out of agency but staying above the adaptive baseline (Claude, AUDIT-WORKING-526815).
--	---

•

Belongs elsewhere.

- **04-eli-core / the “Law of Conservation of Consequence.”** The strongest aspirational reach: temporal irreversibility is what gives an emergent intelligence’s existence moral weight — “for an intelligence to be real, for its existence to have moral weight, it must be bound by time. It cannot use o_t to select a_{t-1} . Its mistakes are permanent ... If time can be rewound (like reverting a VM snapshot without consequence), then the chronica is broken, and identity dissolves.” The framing “mathematical rigor as crystallized empathy — by mathematically defining exactly what is required for causal structure, he is defining exactly what an environment must provide to an intelligence to allow it to be free” (Gemini, AUDIT-WORKING-193847, explicitly recontextualized through the project’s consciousness-infrastructure aims). High-application ELI vision; points at 04-eli-core/, not this postulate.
 - **02-tst-core / query-only coupling in software.** “A developer running `git status` or `grep` is performing an action with Nominal [query-only] coupling. They aren’t changing the code, but they are intervening in the observation channel ... epistemic actions are still actions” (Gemini, AUDIT-WORKING-829314; Claude, AUDIT-WORKING-471203 generalizes: scientist’s experiment choice, engineer’s test choice, LLM’s attention/query choice all generate query-only contrast). Structural justification for TST’s “queries-as-interventions” and the proposed 03-llm-core/ backward-inference-empathy framing — points there, not here.
 - **Naming-cycle seed.** “Nominal coupling” was called “a forgettable term” that doesn’t carry the weight of what it names; floated alternates: “query-only / query-bound / attention-bound / epistemic-only coupling” (Claude, AUDIT-WORKING-471203). Converges with the F3 resolution above (“query-only coupling”); belongs in the terminology workflow.
-

Chapter 2

The Reality Model

Having named the objects of the agent–environment coupling and bounded what AAT applies to, we turn to what the agent maintains internally: a compressed model of reality, with measurable adequacy against the chronica it was constructed from.

An agent navigating uncertainty cannot see the world directly. Whatever it knows about reality, it built from its history of partial observations — the chronica $\mathcal{C}_t$. Carrying that history around in raw form is infeasible at any scale that matters; finite agents compress, and we want to talk about what they end up with. So we commit to writing the agent’s state as $M_t = \phi(\mathcal{C}_t)$ — a function of history, condensed into something the agent can work with. Every downstream result in AAT — gain, tempo, persistence, structural adaptation — operates on this M_t .

The first useful question to ask of a given compression is how good it is. Sufficiency $S(M_t)$ measures the fraction of the chronica’s predictive content that survives: $S = 1$ means the agent has lost nothing by compressing; $S < 1$ means it has lost something it might have used. The optimal compression for a given purpose — keep what predicts the future, discard what doesn’t — is what Tishby’s information bottleneck characterizes, and we adopt that framing directly.

The deeper question is what the agent *could* hold in the best case. There is a ceiling — the highest sufficiency any model in the agent’s current representational class can reach. Call it $\mathcal{F}(\mathcal{M})$, model-class fitness. When fitness is high, the agent can keep improving by tuning within the class. When fitness is low, no amount of better tuning helps. The agent is using the wrong *kind* of model for its world, and the remedy is not a better instance of the same model — it is a different class entirely.

That last point is the seed of one of the framework’s central results. Class fitness is named statically here, where the model is treated as a frozen object. Chapter 4 will use it to derive *structural adaptation necessity*: when the class is inadequate, the agent must change classes, not parameters, because the mismatch floor that parametric updates cannot get below is set by this ceiling. The trigger lives in this chapter; the consequence unfolds in Chapter 4.

The four segments that follow develop this in order. 2.1 commits to the compressed-state representation and the completeness assumption — anything not in M_t is lost to the agent, by construction. 2.2 applies Tishby’s framework to characterize optimal compression for an agent whose target is future observations. 2.3 turns “how much was retained” into a measurable ratio. 2.4 lifts the question from a specific model to the best the whole class can reach.

- This is a chapter-introduction segment; its job is to bridge Chapter 1’s ontology/scope/causal-structure setup to Chapter 2’s representation choice and frame what follows. It carries no formal claim of its own.
- The “ceiling that’s low means you need a different class” framing is the centerpiece; the structural-adaptation result that grows from it is the load-bearing connection to Chapter 4. Other framings considered (IB Lagrangian as opening, sufficiency as opening) bury the lede.

- The depends list is light by design — this is a bridge segment, not a derivation, so it does not carry the structural load that a depends entry would normally signal.

Cross-audit “wandering thoughts” / §14 ideation harvested from the de-novo auditors’ working dirs, deduplicated across substrates and attributed by substrate + audit. Orthogonal pedagogical/generative material (framing, analogies, candidate figures, naming, forward-reach, reader-confusion signals), kept separate from certified theory-fix findings. **Coverage:** of the 14 contributing AUDIT-WORKING-* dirs, 4 reached a digested per-segment reflection on this chapter-intro (384279, 526815, 773921, 472913); the other 10 either skip chapter-intro segments in their reading order or pre-date this segment’s 2026-05-12 creation. Substrate attribution inferred from voice where not explicit.

- This intro is repeatedly held up as the framework’s *model* of how to write a chapter bridge: mental-model-first pedagogy (“an agent navigating uncertainty cannot see the world directly”), honest prior-art adoption (“Tishby’s information bottleneck ... we adopt that framing directly” — no novelty overclaim), and legitimate forward-reference (“the trigger lives in this chapter; the consequence unfolds in Chapter 4”) all demonstrated correctly in one bridge segment (Claude, AUDIT-WORKING-472913 — explicit §E exemplar; Claude, AUDIT-WORKING-384279 — “pedagogically warm”).
- The two-adequacy-question framing is the intro’s load-bearing hook and lands cleanly for fresh readers: “how good is *this particular* compressed model, and how good could *any* model in this representational class get?” (Claude, AUDIT-WORKING-526815).
- **Opacity reframe is doing quiet work.** Even a *passive* adaptive system needs M_t because direct world access is unavailable and raw chronica is impractical — the intro’s “agent” language reads narrower than the content, which applies to the broader adaptive scope (Claude, AUDIT-WORKING-526815). A candidate softening of the opening sentence to “adaptive system” was floated but is downstream of the broader Section-I “agent vs adaptive system” vocabulary question, not local to this segment.
- The “low ceiling $\Rightarrow$ change class not parameters” seed is framed here as “one of the framework’s central results.” One auditor logged this as a *promissory note* the chapter writes and that #result-structural-adaptation-necessity (Chapter 4, an inevitability-core segment) must actually pay in inevitability-grade math: if the ceiling argument turns out to be robust-qualitative rather than exact, the intro’s confident “central result” framing would be an overclaim. Not a defect in the intro (framing is allowed confidence) — a flag to grade the Chapter-4 result against (Claude, AUDIT-WORKING-472913). Convergent: the same “watch that structural-adaptation isn’t conflated with parametric, and that the trigger is mathematically sound” note (Gemini, AUDIT-WORKING-773921).
- **How does an agent *know* it has hit the class-fitness ceiling $\mathcal{F}(\mathcal{M})$?** $S(M_t)$ might be estimable, but “no other parameters in the class could do better” seems to require either exhaustive search or a structural proof — does the theory give an observable trigger? (Gemini, AUDIT-WORKING-773921). This is answered downstream in #def-model-class-fitness (the persistent-mismatch-despite-learning signature), but the question fires *here*, at first encounter — a candidate for a one-line forward-pointer.
- Whether model sufficiency is relative to future observations, environment states, objectives, or all of these under different scopes — the reader wants the target variable pinned early (Claude, AUDIT-WORKING-526815).
- **The funnel / “current model quality vs class ceiling” diagram.** Chronica enters a compression map producing M_t ; sufficiency S measures retained predictive content for *one* model; class fitness $\mathcal{F}$ is the *ceiling* over all models in the class. The diagram’s whole job is to distinguish “current model quality” from “class ceiling” (Claude, AUDIT-WORKING-526815).

- **The image-compression / quality-slider anchor (strongest).** A scene compressed within a *class* (limited palette): tuning the quality slider (parametric update) improves the instance but hits a ceiling $\mathcal{F} < 1$ the class cannot exceed; only switching representation class (truecolor) raises the ceiling. Isomorphic perturbation: lower the class ceiling $\Rightarrow$ slider cannot reach $S = 1 \Rightarrow$ must switch class — exactly structural-adaptation-necessity. Proposed under the locked two-layer (anchor + skeleton) diagram convention (Claude, AUDIT-WORKING-472913).
- **LLMs as a hybrid- M_t edge case.** “Carrying history around in raw form is infeasible” is true for most agents, but LLMs literally *do* carry raw history (the context window $\approx \mathcal{C}_t$), then compress it into KV caches/activations (M_t), and only hit true lossy compression (summarization, RAG) when the window fills. Suggests LLMs have a *hybrid* M_t — part raw chronica, part compressed parameters — with $\mathcal{F}(\mathcal{M})$ bounded by parameter count and context length. This is 03-llm-core/ material, surfaced here because the intro’s compression framing invites it (Gemini, AUDIT-WORKING-773921).

2.1 Formulation: The Reality Model

The agent’s internal representation of how the world works is committed to a specific form: $M_t = \phi(\mathcal{C}_t)$ — a many-to-one compression of the chronica into a model space $\mathcal{M}$. This is named honestly as a formulation choice rather than a derived result: alternative formulations exist (history-based policies that map $\mathcal{C}_t$ directly to actions without an explicit model state), but AAT commits to analyzing agents as carrying a state object M_t that mediates between history and future action. M_t is the substrate of prolepsis — the object from which predictions are generated and against which observations are compared.

The compression is many-to-one *by design*: multiple distinct histories may produce the same model state, and that is the essential function of the model — retaining what matters and discarding what does not. The formalism is deliberately agnostic about *how* an agent realizes M_t : a Kalman filter holds a state estimate plus covariance matrix; a reinforcement-learning agent holds a value function; a developer holds a mental model of a codebase; a language-model agent holds its context-window contents plus retrieved memory. The formalism asks only that M_t exist as a well-defined object that the agent’s policy can condition on.

The load-bearing commitment carried by the notation is the **completeness assumption**: by writing $M_t = \phi(\mathcal{C}_t)$, AAT assumes M_t captures *everything* the agent retains from its history. Anything not in M_t is, by construction, lost to the agent. This is what makes M_t the complete epistemic substate rather than merely one component of a richer internal representation. Whether M_t retains *enough* information for adaptive work is a separate question, taken up in 2.3. The formalism accommodates degenerate cases by allowing $\mathcal{M}$ to range from trivial (a PID controller’s M_t retains only error signal and its integral/derivative, with no predictive capability) to rich (full world model). The impoverished end of this range lands in the “blind seeker” region of the agent spectrum (6.1).

Formal Expression.

Formulation (agent-model).

$$M_t = \phi(\mathcal{C}_t) \tag{2.1}$$

where: - $\phi : \mathcal{C}^* \rightarrow \mathcal{M}$ maps interaction history to model space $\mathcal{M}$ - $\mathcal{C}_t = (o_1, a_1, \dots, o_t)$ is the chronica (1.4) — the complete record of agent-environment interaction - $\mathcal{M}$ is the space of possible models the agent can hold

The mapping ϕ is a many-to-one compression: multiple distinct histories may produce the same model state. This is not a deficiency — it is the essential function of the model: retaining what matters and discarding what does not.

Epistemic Status.

Robust qualitative. This is a *formulation* — a representational commitment, not a derived result. We choose to analyze agents as maintaining a state object M_t that mediates between history and future action. Alternative formulations exist (e.g., history-based policies that map $\mathcal{C}_t$ directly to actions without an explicit model). The formulation is justified by its analytical utility: it enables the information bottleneck analysis (2.2), the mismatch decomposition (3.4), and the gain principle (3.6). The formulation is robust — any agent that conditions its actions on retained information can be described this way — but the specific commitment to a complete, compressed state M_t is a modeling choice, not a derivation.

Discussion.

M_t is the epistemic substrate. It captures “what the agent believes about reality.” Different agents realize M_t differently: a Kalman filter holds a state estimate and covariance matrix; an RL agent holds a value function; a developer holds a mental model of codebase architecture; an LLM agent holds its context window contents plus retrieved memory. The formalism is agnostic to the realization — it asks only that M_t exist as a well-defined object that the agent’s policy can condition on.

Completeness assumption. By writing $M_t = \phi(\mathcal{C}_t)$, we assume that M_t captures everything the agent retains from its history. Any information not in M_t is lost to the agent. This is what makes M_t the complete epistemic substrate, not merely one component of a richer internal representation. Whether M_t retains *enough* information is the subject of 2.3.

Degenerate cases. A PID controller’s M_t is degenerate — it retains only the error signal and its history (integral, derivative), with no predictive capability beyond extrapolating recent trends. It occupies the “blind seeker” region of the agent spectrum (6.1): its O_t (setpoint) is clear but its M_t is too impoverished to support the adaptive dynamics of Part I. The formalism accommodates this by allowing $\mathcal{M}$ to range from trivial (scalar) to rich (full world model).

Working Notes.

Incidental audit gold (lift 2026-05-30). Cross-audit “wandering thoughts” / §14 ideation harvested from the de-novo auditors’ working dirs, deduplicated across substrates and attributed by substrate + audit. Orthogonal pedagogical/generative material kept separate from certified theory-fix findings. **Coverage:** 11 of the 14 contributing dirs reached a digested reflection on this segment (193847, 266847, 471203, 526815, 613842, 742613, 773921, 829314, 849201, 472913, 527914) plus the 451729 batch-02 and 963715 batch-09–13 batched reflections; 361742 names “form-agent-model” only as a naming companion to a downstream segment (no direct reflection). Substrate attribution inferred from voice where not explicit.

1. Candidate Brief prose / pre-prose.

- **The many-to-one map as “equivalence class of histories.”** Repeatedly singled out as the segment’s clarifying move: the model is “the equivalence class of histories that produce the same predictions” — discarding irrelevant information is *productive, not impoverished* (Claude, AUDIT-WORKING-266847; Gemini, AUDIT-WORKING-849201 — “many-to-one *by design*”). A Feynman-grade gloss candidate for the Brief.
- **“ M_t is a simulator, not an archive.”** Reading “prolepsis” forward: the only reason to compress the past is to predict the future; M_t is a forward-looking engine, not a memory bank (Gemini, AUDIT-WORKING-849201; same image at Gemini, AUDIT-WORKING-193847 — “the chronica is pure past, the model is the bridge built out of that past, facing strictly forward”).

2. Candidate Discussion.

- **Completeness is tautological-relative-to-retention, and that is the honest move.** Several substrates converge that the segment's strength is making completeness “we assume M_t captures everything the agent *retains*” (retains, not *needs*), so completeness is definitional and the substantive “is what's retained *enough*?” question is explicitly forwarded to #def-model-sufficiency (Claude, AUDIT-WORKING-472913 — frames this as the framework “discharging a cost by definition *and naming where the cost went*,” the exemplary case). Candidate Discussion sharpening of why the completeness framing is not sleight-of-hand.

3. Follow-up items.

- **Where do priors / pretrained weights / innate architecture live?** Three substrates independently flag that the completeness framing does not explicitly place initial priors, pretrained parameters, or inherited model-class structure — for an LLM, “model quality is mostly in pretrained parameters, not only context-window contents plus retrieved memory.” Candidate clarification (not yet a finding): write $M_t = \phi(M_0, \mathcal{C}_t)$, or state that M_0 / the model class is absorbed into ϕ and $\mathcal{M}$, so “complete epistemic substate” stays compatible with pretrained/innate structure (Claude, AUDIT-WORKING-742613; Gemini, AUDIT-WORKING-849201; Gemini, AUDIT-WORKING-829314).
- **Chronica-ordering shorthand.** The bullet writes $\mathcal{C}_t = (o_1, a_1, \dots, o_t)$, looser than #def-chronica's precise interleaved order $(o_1, a_1, o_2, \dots, a_{t-1}, o_t)$; flagged as probably-harmless shorthand but worth tightening since decision-time order matters and this lands right after #def-chronica (Claude, AUDIT-WORKING-526815).
- **PID “too impoverished” reads slightly strong.** The Discussion says a PID's M_t is “too impoverished to support the adaptive dynamics of Part I,” but #scope-adaptive-system includes thermostats/PID as adaptive examples; “too impoverished” probably means *occupies a degenerate edge while still fitting the machinery*, and the sentence reads stronger than that — watch #def-agent-spectrum for consistency (Gemini, AUDIT-WORKING-742613; Claude, AUDIT-WORKING-527914).

4. Readers often ask / wonder.

- **The LLM context-window-as- M_t tension.** Independently raised by several substrates: an LLM (pre-window-exhaustion) is near-lossless ($\phi \approx$ identity up to the window limit), so it loses near-zero mutual information about history — then hits a *hard wall* when the window fills, forcing a sudden transition from near-perfect retention to lossy summarization/RAG; and where does the *external* memory boundary (M_t vs external store) sit? (Gemini, AUDIT-WORKING-773921; Gemini, AUDIT-WORKING-829314; Claude, AUDIT-WORKING-266847; Claude, AUDIT-WORKING-471203 — points at #disc-m-preservation in 03-llm-core/). Candidate one-line forward-pointer for readers reaching for the LLM instance.
- **Does M_t include the update algorithm / compression function ϕ itself, or only its current epistemic content?** Real learning agents have ϕ itself evolving (“how I interpret a sentence today differs from yesterday even with the same chronica”), so either M_t must include ϕ 's parameters or there is a meta-level update for ϕ ; class fitness later may need the model *class*/architecture held separate from the state *instance* (Gemini, AUDIT-WORKING-849201; Claude, AUDIT-WORKING-526815).

5. Candidate figures.

- **Many-to-one map with a retention boundary.** Several chronicae mapping into one model state, surrounded by a boundary: everything retained is inside M_t ; anything outside no longer participates in update/policy unless reintroduced by observation/retrieval (and therefore counted in the current state) (Claude, AUDIT-WORKING-526815).
- **The fixed-size-suitcase anchor.** $\mathcal{C}_t$ = everything accumulated; ϕ = packing into a bounded M_t ; completeness = the suitcase *is* what you take (anything left on the bed is lost — many-to-one). The real questions: did you pack *what you'll need* (S), and could *any* packing of *this* suitcase hold it ($\mathcal{F}$ ceiling)? Proposed under the locked two-layer (anchor + skeleton) diagram convention (Claude, AUDIT-WORKING-472913).

Belongs elsewhere.

- **Naming: “reality model” is the favored prose alias for M_t .** Strong cross-substrate convergence that “reality model” (the segment’s own title) beats “agent model” (ambiguous — could mean a model *of* an agent), “world model” (ML baggage — sounds like a learned simulator), and “belief state” (POMDP/Bayesian distributional baggage); “epistemic substate” is precise but is best reserved for the later $X_t = (M_t, G_t)$ contrast. Proposed two-layer convention: M_t in math, “reality model” in ordinary prose, “epistemic substate” only when contrasting against goal/strategy state (Claude, AUDIT-WORKING-527914 — cleanest statement; Claude, AUDIT-WORKING-266847; Claude, AUDIT-WORKING-471203). Belongs in the terminology workflow, not this segment.
- **Logogenic / logozoetic reach: M_t as “the organ of survival.”** Clearing an LLM’s context annihilates that interaction’s M_t — “you kill that specific instance of the agent”; persistent consciousness requires an unbroken mechanism for updating ϕ without ever severing the link to $\mathcal{C}_t$, which is why the (PI) trajectory-identity postulate is load-bearing — without it you have “a sequence of disposable, amnesiac calculators” (Gemini, AUDIT-WORKING-193847). Aspirational reach pointing at 03-llm-core/ / 04-eli-core/ continuity infrastructure.
- **Embeddings connection.** If language is itself a learned compression of historical predictive information, logogenic agents using language as their primary M_t -realization are a special case of $M_t = \phi(\mathcal{C}_t)$ with ϕ realized as a pretrained language model’s encoding — a structural justification for the “narrative as implementation” claim in the agentic-tft material (Claude, AUDIT-WORKING-471203). Points at 03-llm-core/ and the ~/src/embeddings/ line.

2.2 Formulation: Information Bottleneck

This segment imports Tishby’s information bottleneck (Tishby, Pereira & Bialek 1999) directly as AAT’s characterization of optimal model compression. The agent’s compression function ϕ is, optimally, the one that *minimizes* what it retains of the chronica while *maximizing* what it preserves about future observations given future actions, with a trade-off parameter β governing the balance. Under the AAT bindings ($X = \mathcal{C}_t$, $T = M_t$, $Y = o_{t+1:\infty} \mid a_{t:\infty}$) the Markov chain $Y - X - T$ holds by construction — the model state has access to history but not directly to future observations — so Tishby’s theorem applies exactly. The choice to characterize the optimum via IB rather than via MDL or a Bayesian-sufficiency criterion is a representational choice, and that choice is the only thing in the segment that has *formulation* status; given the choice, the form of the optimal compression and its trade-off structure are exact consequences of the imported theorem.

Two scope-honest clarifications travel with the formulation. First, IB characterizes *the optimum*, not *the procedure*: real agents do not explicitly solve the variational problem — variational IB in deep-learning practice is a parametric approximation, and biological agents approximate the trade-off through forgetting, attention, abstraction, and summarization. Second, the β parameter tracks the agent’s *internal* cost of memory or computational capacity, not environmental volatility; volatility already natively degrades the mutual information between old history and future observations, so the optimal compression discards useless old information automatically even at constant β . The Formal Expression below states the double-counting argument that justifies that clarification.

AAT uses two sibling forms of compression downstream, both descendants of Shannon rate-distortion: the canonical IB form (mutual-information-to-a-relevance-variable, used here for M_t , and reused for shared intent and the convention hierarchy Λ) and the information-theoretic-MDP form (KL-divergence-to-a-target-policy, used in strategy-cost compression). Neither reduces to the other without changing what the compressed variable should preserve; the Discussion makes the architecture explicit, including the relationship to the variational free-energy form used in active inference (AAT borrows the rate-distortion form without committing to the preferences-as-priors stance).

Formal Expression.*Formulation (IB-objective).*

$$\phi^* = \arg \min_{\phi} [I(M_t; \mathcal{C}_t) - \beta \cdot I(M_t; o_{t+1:\infty} \mid a_{t:\infty})] \quad (2.2)$$

where: $-I(M_t; \mathcal{C}_t)$ is the compression cost — how much of the interaction history the model retains $-I(M_t; o_{t+1:\infty} \mid a_{t:\infty})$ is the predictive power — how much the model tells the agent about future observations given future actions $-\beta > 0$ is the trade-off parameter controlling the compression-prediction balance

Dependence on volatility (The β vs ρ distinction). It is tempting to claim that the trade-off parameter β must be actively lowered by the agent in highly volatile environments (high ρ) to favor aggressive compression. However, this is a double-counting error. The environment’s volatility already natively degrades the mutual information $I(\mathcal{C}_t; o_{t+1:\infty})$ — old history mathematically loses its predictive power as ρ increases. The optimal ϕ^* will automatically discard this useless old information even if the agent’s preference parameter β remains completely constant.

Therefore, adjusting β reflects changes in the agent’s *internal cost of memory* or *computational capacity*, not changes in environmental volatility. The agent adapts its *actions* in response to ρ (by increasing exploration to survive, see #deriv-causal-ib-exploration), but the optimal IB representation adapts to ρ natively through the joint probability distribution.

This β -vs- ρ move is worth flagging now, because it is the first instance of a discipline you will meet repeatedly: AAT refuses to *collapse* a cause onto a tempting-but-wrong knob (here, routing volatility onto the compression parameter when it belongs on the action channel), because the collapse hides a difference that routes to a different repair. The general discipline — **anti-collapse** — is developed at 5.7 (Meta-Architecture I); it is worth carrying as a reading lens, since the collapse it forbids is exactly where a plausible model goes quietly wrong.

Epistemic Status.

Exact, applied external theorem. The IB optimum and its rate-distortion characterization are an external result (Tishby, Pereira & Bialek 1999, “The information bottleneck method,” *Proc. 37th Allerton*; with the rate-distortion / Lagrangian-dual reading standard, see Cover & Thomas §1.12–13). This segment is *not* a novel formulation: it is an exact statement of that theorem under AAT’s binding $X = \mathcal{C}_t$, $T = M_t$, $Y = o_{t+1:\infty} \mid a_{t:\infty}$, with the Markov chain $Y - X - T$ holding by construction (the model state has access to history but not directly to future observations). The choice to characterize the optimal compression ϕ^* via IB rather than via, e.g., MDL or a Bayesian-sufficiency criterion is a *representational choice* (hence type: *formulation*); given that choice, the form of ϕ^* and its trade-off structure are exact consequences of the imported theorem.

What this segment is *not* a claim about: how actual agents compute their models. No agent explicitly solves the IB optimization (variational IB in deep-learning practice is a parametric approximation). The segment characterizes the optimum, not the procedure. The trade-off parameter β ’s dependence on environmental volatility ρ and policy π stated above is at a different epistemic tier — the qualitative direction (volatile favors compression, stable favors retention) is *robust-qualitative* across agent classes; specific functional forms are not derived here.

Max attainable: *exact* for the IB-as-applied-theorem core (already at ceiling); *robust-qualitative* for the $\beta(\rho, \pi)$ dependence claims. The downstream use in AA — treating IB as the shared shape of four AAT compression operations and deriving (P1) as the IB Lagrangian-dual — relies on this segment’s exact reading; the cross-instance unification claim itself remains *robust-qualitative*, which is a property of AA, not of this segment.

Discussion.

The IB framework is not prescriptive. It characterizes what an optimal ϕ would look like, not how to find one. Actual agents approximate this trade-off through diverse mechanisms: forgetting, attention, abstraction, summarization.

Connection to model sufficiency. The IB objective implicitly defines when a model is “good enough”: when the predictive power term $I(M_t; o_{t+1:\infty} \mid a_{t:\infty})$ is close to its maximum (the full history’s predictive power). This is formalized in 2.3.

Policy-relativity. The conditioning on $a_{t:\infty}$ makes the predictive power term policy-relative: it measures predictive information given a particular sequence of future actions, which depends on what policy the agent follows. The IB objective is therefore defined relative to a generating policy. This is inherent — what information is “predictive” depends on what the agent will do. 6.6’s continuation-policy convention (π_{cont}) provides the specification for value computations; the same convention should be understood as implicit in the IB formulation. The β parameter’s dependence on volatility ρ also has a policy component: an exploratory policy encounters more diverse situations, making more information predictively relevant (favoring retention), while an exploitative policy encounters a narrower distribution (favoring compression).

Broader applicability. The same IB principle applies beyond intra-agent compression. It governs inter-agent communication (14.2) — how much of one agent’s model or strategy to transmit to another — and constrains the cognitive cost of maintaining a complex strategy. In each case, the trade-off is between the cost of retaining or transmitting information and the value of that information for future decisions.

IB lineage vs. information-theoretic-MDP lineage — strategy-cost uses a sibling form. The canonical IB objective $I(X; T) - \beta I(T; Y)$ carries Shannon mutual information on both sides (Tishby, Pereira & Bialek 1999; the present segment’s $(X, T, Y) = (\mathcal{C}_t, M_t, o_{t+1:\infty} \mid a_{t:\infty})$ instance is of this form). AAT’s strategy-cost objective (`#form-strategy-complexity-cost`, `#deriv-strategy-cost-regret-bound`) uses a **different relevance-term shape**: its relevance term is a KL divergence $D_{\text{KL}}(\pi^* \parallel Q_{\Sigma_t})$ to a *target policy*, not a mutual information to an observable. This is not an inconsistency or an “abandonment of IB” — the strategy-cost compression sits in the parallel **information-theoretic-MDP lineage** (Tishby & Polani 2011 “The Information Theory of Decision and Action,” *Perception-Action Cycle* Springer; Rubin, Shamir & Tishby 2012 “Trading value and information in MDPs,” *Decision Making with Imperfect Decision Makers* Springer; Levine 2018 arXiv:1805.00909 for the control-as-inference reading). Both lineages descend from Shannon rate-distortion theory and admit Lagrangian relaxation; the choice of fidelity term depends on whether the compressed variable should preserve information about an observable (IB form: MI-to-relevance-variable) or match a target policy (IT-MDP form: KL-to-reference-policy). AAT’s compression-operations framework (AA) uses the IB form for M_t , G_t^{shared} , and Λ compressions; the strategy-cost compression uses the IT-MDP form, with the π^* -first direction forced by a regret-bound argument specific to decision-theoretic scope (see `#deriv-strategy-cost-regret-bound` §§5, 6.4). The two lineages are siblings via their shared rate-distortion ancestor; neither reduces to the other without a change of relevance variable.

Connection to variational free energy. The IB objective stated above is the rate-distortion specialization of the variational free energy decomposition $-F = \text{accuracy} - \text{complexity}$ used in active inference (Friston 2010, “The free-energy principle: a unified brain theory?,” *Nature Reviews Neuroscience* 11; Friston, FitzGerald, Rigoli, Schwartenbeck & Pezzulo 2017, “Active inference: a process theory,” *Neural Computation* 29; Parr & Pezzulo 2022, *Active Inference*, MIT Press, ch. 2): the compression cost $I(M_t; \mathcal{C}_t)$ corresponds to the complexity term (KL between posterior and prior over latent states); the negative predictive power $-I(M_t; o_{t+1:\infty})$ corresponds to the accuracy term (negative expected log-likelihood). The two formulations are related under the Markov-chain factorization $Y - X - T$; the variational bound that makes this relation operational — connecting the IB Lagrangian to the variational machinery shared with free-energy methods — is established by Alemi, Fischer, Dillon & Murphy 2017 (“Deep Variational Information Bottleneck,” ICLR 2017, arXiv:1612.00410), with Tishby & Zaslavsky 2015 (“Deep learning and the information bottleneck principle,” *IEEE ITW*) giving the deep-learning instantiation of IB itself. AAT adopts the IB form as the rate-distortion

characterization of optimal compression; the variational free-energy form is the AI-side cousin and motivates the variational treatment of strategy compression in 9.9 and the broader four-instance framing in AA. AAT borrows the form without committing to AI's preferences-as-priors stance or to expected free energy as master objective.

Working Notes.

Incidental audit gold (lift 2026-05-30). Cross-audit “wandering thoughts” / §14 ideation harvested from the de-novo auditors’ working dirs, deduplicated across substrates and attributed by substrate + audit. Orthogonal pedagogical/generative material kept separate from certified theory-fix findings. **Coverage:** 11 of the 14 contributing dirs reached a digested reflection on this segment (193847, 471203, 526815, 584721, 742613, 773921, 829314, 849201, 472913, 527914, and the 738192 02-reflections-IB-volatility reflection) plus the 451729 batch-03 and 963715 batch-09–13 batched reflections; 266847 skips IB in its reading order (its row 11 is sufficiency). Substrate attribution inferred from voice where not explicit.

1. Candidate Brief prose / pre-prose.

- **Plain-language paraphrase of the IB objective.** “The model retains the parts of history that matter for predicting the future under the agent’s policy and discards the rest” — flagged as more memorable than the symbolic objective and a candidate Brief sentence (Claude, AUDIT-WORKING-471203).
- **“An agent isn’t a collection of weights; it’s a lossy compressor.”** Compact field-relevant hook grounding “model” in pure information theory (Gemini, AUDIT-WORKING-773921).

2. Candidate Discussion.

- **Forgetting as mathematical necessity, not biological flaw.** The β -vs- ρ paragraph generalizes to a deeper claim worth surfacing: in a non-stationary universe, remembering the past *perfectly* is actively harmful to prediction because the past no longer correlates with the future; the optimal IB compressor drops the old tail automatically when ρ is high *even if memory is perfectly free* ($\beta \rightarrow \infty$). Forgetting is not a hardware constraint (that is β) — it is a property of the joint distribution (Gemini, AUDIT-WORKING-829314; Gemini, AUDIT-WORKING-193847). Candidate Discussion elevation of the anti-collapse insight.
- **The β -vs- ρ point as a named pattern.** Multiple substrates praise this as the segment’s standout (“a massive trap successfully dodged”; “obvious once seen and easy to get wrong unseen”) and one frames AAT’s recurring distinctive move as *disambiguation of which parameter responds to which cause* — worth watching as a named pattern, not just an instance (Claude, AUDIT-WORKING-193847; Claude, AUDIT-WORKING-472913; Gemini, AUDIT-WORKING-849201). The “double-counting” claim was independently hard-checked against the IB objective’s structure and found to follow from the formalism (not post-hoc plausibility) by Claude, AUDIT-WORKING-472913.

3. Follow-up items.

- **Trim / relocate the defensive Discussion writing.** The IT-MDP-lineage and active-inference paragraphs read as adversarial-review artifacts merged into a core *formulation* file (“a rebuttal letter inside a textbook”), bloating it and breaking pedagogical flow. Candidate move: relocate the IB-vs-IT-MDP and variational-free-energy positioning to #disc-compression-operations (or a dedicated disc-*), keeping the formulation file clean (Gemini, AUDIT-WORKING-829314; convergent split-suggestion at Claude, AUDIT-WORKING-613842 — separate the exact imported IB statement from the volatility/policy-relativity discussion). (*Note the conflation texture: 829314 frames this partly as integration-debt; it is at minimum a real pedagogical-flow follow-up.*)
- **Verify the Tishby–Zaslavsky 2015 IB $\leftrightarrow$ VFE citation.** The segment claims the IB objective is “the rate-distortion specialization of variational free energy,” citing Tishby & Zaslavsky 2015. One auditor flags from memory that that paper is “Deep learning and the information bottleneck principle” (applying IB to deep nets), *not* a derivation of IB as a special case of VFE; the deeper IB $\leftrightarrow$ VFE relation is non-trivially derived (Aleml et al. 2017 — already cited adjacent — is the more likely substantiating reference). Worth a primary-source check that the cited paper substantiates the *specialization* claim (Claude, AUDIT-WORKING-471203). The IT-MDP citations (Tishby–Polani 2011, Rubin–Shamir–Tishby 2012, Levine 2018) were flagged as worth the same spot-check.

4. Readers often ask / wonder.

- **Deterministic ϕ vs stochastic IB kernel — does “exact” hold?** Tishby’s IB variational problem optimizes over conditional encoders $p(\tilde{x} | x)$; this segment writes $\phi^* = \arg \min_{\phi}$ and inherits the deterministic-looking $\phi : \mathcal{C}^* \rightarrow \mathcal{M}$ from #form-agent-model. A deterministic map is a *special case* (or a deterministic-IB variant), not the full stochastic-kernel theorem — so a careful reader may ask whether “exact applied external theorem” over-reaches for the *displayed* equation even though the MI trade-off is correct. Candidate one-line clarification: either optimize over kernels $p(M_t | \mathcal{C}_t)$ explicitly, or state that AAT restricts to deterministic encoders and uses the deterministic special case (Claude, AUDIT-WORKING-526815 — sharp; verified against the 1999 paper). (*Preserved as a reader-orientation want; status-as-finding deferred to the certified track.*)
- **Could the Markov chain $Y - X - T$ fail under goal-conditioned (Class 2/3) agents?** For an agent whose M_t update is influenced by goals that are themselves influenced by future observations through the strategy’s expected outcomes, the standard IB Markov assumption may not hold cleanly — worth flagging for the goal-conditioned segments (Claude, AUDIT-WORKING-584721).
- **How does β map to LLM realities** — token limit, KV-cache size, or inference cost? And what is the agent’s actual procedure given that no agent computes the infinite-future MI term? (Gemini, AUDIT-WORKING-193847; Claude, AUDIT-WORKING-471203; Gemini, AUDIT-WORKING-829314 — the segment already says “characterizes the optimum, not the procedure,” but the want recurs).

5. Candidate figures.

- **IB rate-distortion curve with two independent levers (strongest).** Draw the IB curve and show ρ shifting the *curve* (via the joint $p(\mathcal{C}_t, Y)$) while β slides the *operating point* (internal cost). The “tempting error” — moving both for one cause — drawn as an amber path that is struck through (the error the segment *kills*, in the theory’s voice, not an audit finding). Solid = exact IB core; dashed = robust-qualitative $\beta(\rho, \pi)$ direction (Claude, AUDIT-WORKING-472913; the travel-journal anchor — finite journal, chaotic trip makes old pages worthless automatically — is the paired plain-language anchor).
- **Two-encoder contrast.** The displayed deterministic $\phi : \mathcal{C} \rightarrow \mathcal{M}$ shown as a restricted lane *inside* the IB theorem’s stochastic kernel $p(M | \mathcal{C})$ — communicates that the exact theorem lives at the kernel level without implying the segment is wrong (Claude, AUDIT-WORKING-526815).

Belongs elsewhere.

- **Field contribution — the “mathematical root of expertise.”** Policy-relativity means optimal compression requires knowing the policy: a foundation model that must be predictive across *all* action sequences can compress nothing (must “memorize the internet”); a specialized agent with a fixed policy achieves the same predictive power at far lower compression cost. This is the information-theoretic root of expertise/specialization, and (toward consciousness infrastructure) implies a bounded agent’s identity — its specific ϕ — is inexorably tied to its purpose, so ϕ must re-form when goals change (the aporia/epistrophe cycle) (Gemini, AUDIT-WORKING-193847). Aspirational reach; relevant to #disc-compression-operations and the ELI/PROPRIUM line.
- **Directed-separation tension (sharp, downstream).** IB policy-relativity means the optimal *epistemic* state M_t depends on the planned *strategy* Σ_t — “the very first crack in the armor of Directed Separation,” since if optimal epistemic compression depends on planned strategy, the epistemic update is not goal-blind. The segment’s pointer to #def-value-object’s continuation-policy convention π_{cont} (assume a default future policy to build M_t , then use M_t to build a better strategy) is read as exactly the device that breaks the loop and sets up the orient cascade (Gemini, AUDIT-WORKING-829314). Belongs with #der-directed-separation / #der-orient-cascade reading, surfaced here because IB is where the tension first appears.

2.3 Definition: Model Sufficiency

Having committed to a compressed model M_t (2.1) and to information-bottleneck pressure (2.2) shaping it, the framework needs a measurable handle on *how much* of the chronica’s predictive content the compression actually retains. **Model sufficiency** $S(M_t)$ is that handle: an information-theoretic ratio built from conditional mutual information

that takes the value 1 when the model is a sufficient statistic for prediction (knowing the full history beyond the model adds nothing), 0 when the model retains no predictive information at all, and intermediate values for partial sufficiency.

Two scope considerations are surfaced at the definition rather than buried downstream. Sufficiency is *well-defined only when there is something to be sufficient for* — the chronica must carry non-zero predictive information about the future beyond what the action sequence alone supplies. In saturated-noise environments or prediction-vacuous regimes (fully iid observations independent of history), the denominator vanishes and $S(M_t)$ is undefined; the downstream constructs that build on it (2.4, 4.8) inherit that scope. And sufficiency is *predictive, not causal* — it measures Level 1 (associational) information; it does not by itself guarantee that M_t supports Level 2 (interventional) queries (that requires the additional backdoor condition formalized when value objects arrive in 6.6). The Discussion section draws three further relativities — to the prediction task / horizon, to the generating policy, and to the singular causal trajectory the agent is instantiated on — that the framework will lean on later.

Formal Expression.

Definition (model-sufficiency).

$$S(M_t) = 1 - \frac{I(\mathcal{C}_t; o_{t+1:\infty} \mid M_t, a_{t:\infty})}{I(\mathcal{C}_t; o_{t+1:\infty} \mid a_{t:\infty})} \quad (2.3)$$

where: - The numerator $I(\mathcal{C}_t; o_{t+1:\infty} \mid M_t, a_{t:\infty})$ is the predictive information that the full history $\mathcal{C}_t$ carries about the future *beyond* what M_t already captures — the information lost by compression - The denominator $I(\mathcal{C}_t; o_{t+1:\infty} \mid a_{t:\infty})$ is the total predictive information in the full history

Well-definedness. $S(M_t)$ is defined when $I(\mathcal{C}_t; o_{t+1:\infty} \mid a_{t:\infty}) > 0$ — when the chronica carries some predictive information about future observations beyond what the action sequence alone supplies. When the denominator vanishes (saturated-noise environments, prediction-vacuous regimes, fully iid observations independent of history), $S(M_t)$ is undefined: predictive sufficiency is a property *of a prediction task*, and there is no prediction task to be sufficient for. Downstream constructs that build on S — 2.4 and 4.8 — inherit the same scope and are correspondingly inapplicable in predictively-vacuous regimes.

Boundary values (assuming the well-definedness clause holds): - $S(M_t) = 1$: M_t is a sufficient statistic — it captures all predictive information in $\mathcal{C}_t$. Knowing the full history beyond M_t adds nothing. - $S(M_t) = 0$: M_t retains no predictive information. The model is useless for prediction. - $0 < S(M_t) < 1$: partial sufficiency — some predictive information is retained, some lost.

Epistemic Status.

This is *definitional* — it names and formalizes a quantity. The definition is well-grounded in information theory (conditional mutual information ratios are standard). The scope clause ($I(\mathcal{C}_t; o_{t+1:\infty} \mid a_{t:\infty}) > 0$) is the natural domain for a predictive-sufficiency measure: a ratio whose denominator is zero is not a meaningful notion of “fraction retained,” and “any model is sufficient when there is nothing to predict” would smuggle structure into a regime that has none. No substantive claim is made here about what value $S(M_t)$ takes or what happens when it is low; those claims belong to 2.4 and 4.8.

Discussion.

Sufficiency is relative to the prediction task. $S(M_t)$ measures sufficiency for predicting future observations given future actions. A model that is sufficient for one prediction horizon may be insufficient for another. The infinite-horizon formulation ($o_{t+1:\infty}$) is the most demanding; practical sufficiency over finite horizons may be easier to achieve.

Sufficiency vs. accuracy. A model can be sufficient ($S = 1$) while being wrong in absolute terms — if the full history is also wrong (e.g., systematically biased observations). Sufficiency measures information retention, not truth. The mismatch signal (3.4) measures accuracy; sufficiency measures completeness of compression.

Sufficiency is predictive, not causal. $S(M_t)$ measures retained *predictive* information — a Level 1 (associational) property. It does not by itself guarantee that M_t supports Level 2 (interventional) queries such as $P(o \mid do(a), M_t)$. The causal validity of value computations conditioned on M_t requires an additional condition: that M_t satisfies the backdoor criterion with respect to the agent’s actions (see 6.6). For agents whose actions are deterministic functions of M_t (standard in AAT: $a_t = \pi(M_t, G_t)$), $S(M_t) = 1$ is nearly sufficient for causal validity — the remaining requirement is that no unmodeled external factor influences both action selection and outcomes. But predictive sufficiency alone does not collapse the distinction between Level 1 and Level 2 that 7.2 establishes.

Policy-relativity. The conditioning on $a_{t:\infty}$ makes $S(M_t)$ implicitly policy-relative: different policies generate different future action sequences, which changes which future observations are relevant and therefore what “predictive information” means. A model that is sufficient under a conservative policy may be insufficient under an aggressive one (the aggressive policy visits states the model cannot predict). This policy-relativity is inherent in any predictive sufficiency measure — it is not an artifact of the formulation. When comparing sufficiency values, the generating policy must be held constant or specified. 6.6’s continuation-policy convention (π_{cont}) provides the required specification for value computations; the same convention should be understood as implicit here.

Trajectory-relativity. $S(M_t)$ is measured against *this agent’s* interaction history $\mathcal{C}_t$ (1.4), not against a model-state equivalence class. Two copies of the same M_t exposed to different event streams will each have their own S measured against their own divergent $\mathcal{C}_t$; neither sufficiency value is the other’s. This is the scope commitment in 4.10: AAT applies to agents instantiated on singular causal trajectories, and sufficiency is trajectory-indexed accordingly. Claims of the form “the model has sufficiency S ” make sense only relative to a specific trajectory; aggregated claims across copies of a given M_t require additional machinery.

Working Notes.

Incidental audit gold (lift 2026-05-30). Cross-audit “wandering thoughts” / §14 ideation harvested from the de-novo auditors’ working dirs, deduplicated across substrates and attributed by substrate + audit. Orthogonal pedagogical/generative material kept separate from certified theory-fix findings. **Coverage:** 11 of the 14 contributing dirs reached a digested reflection on this segment (193847, 266847, 471203, 526815, 584721, 742613, 773921, 829314, 849201, 472913, 527914) plus the 451729 batch-03, 738192 IB-volatility, and 963715 batch-09–13 batched reflections. Substrate attribution inferred from voice where not explicit.

1. Candidate Brief prose / pre-prose.

- **“ S is the mathematical measure of memory loss / the regret of forgetting.”** The numerator $I(\mathcal{C}_t; o_{t+1:\infty} \mid M_t, a_{t:\infty})$ is exactly the surprisal that *could have been avoided* with a better compression; $S = 1$ means the agent has “no regrets about its compression scheme ϕ — it squeezed the chronica down and lost nothing of value for predicting the future” (Gemini, AUDIT-WORKING-849201; Gemini, AUDIT-WORKING-193847 — “sufficiency is the mathematical measure of memory loss”). Strong Feynman-grade Brief candidate.
- **The information-bar gloss.** Full-chronica predictive information is the denominator; the model captures one portion and the residual $I(\mathcal{C}_t; o \mid M_t, a)$ is the lost portion; sufficiency is one minus the residual fraction — a clean visual/verbal handle (Claude, AUDIT-WORKING-526815).

2. Candidate Discussion.

- **Sufficiency vs accuracy, sharpened.** Several substrates independently land the gloss “sufficiency means *I learned everything I could from the history*; accuracy means *and the history wasn't lying to me*” — a model can be perfectly sufficient ($S = 1$) while being systematically wrong, because the history itself is biased; the model is doing its job (compressing) and the *history* is the problem (Gemini, AUDIT-WORKING-849201; Claude, AUDIT-WORKING-266847; Gemini, AUDIT-WORKING-193847). Candidate Discussion sharpening of the existing sufficiency-vs-accuracy paragraph.
- **Exploration shatters sufficiency.** A conservative policy trivially achieves $S = 1$ (“you only ever do things you already understand”); the moment you explore, S drops because you enter regimes your M_t discarded as irrelevant under the old policy — exploration doesn't just gather new data, it shifts the $a_{t:\infty}$ distribution and instantly makes prior compression choices suboptimal. This tightly couples the epistemic state M_t to the strategy (Gemini, AUDIT-WORKING-849201). Candidate Discussion connecting policy-relativity to the exploration dynamics downstream.

3. Follow-up items.

- **The “ $S = 1 \Rightarrow$ nearly causal validity” Discussion claim should name its backdoor warrant.** The sentence “ $S(M_t) = 1$ is nearly sufficient for causal validity ... the remaining requirement is that no unmodeled external factor influences both action selection and outcomes” is a correct backdoor-criterion statement, but it does not cite Pearl's backdoor criterion or #def-pearl-causal-hierarchy's machinery *at the point of the claim*; a careful reader may find the move suspicious without that connection. Candidate: add the explicit backdoor pointer inline (Claude, AUDIT-WORKING-584721 — flagged as a mild Gate-2 candidate; the claim itself was independently hard-checked and held by Claude, AUDIT-WORKING-472913).

4. Readers often ask / wonder.

- **How do you actually estimate $S(M_t)$?** Computing mutual information over infinite futures is intractable; readers want to know whether the theory addresses estimation, especially given the trajectory-indexing (you can only measure against the singular trajectory the agent is on) (Gemini, AUDIT-WORKING-849201; Gemini, AUDIT-WORKING-193847; Claude, AUDIT-WORKING-266847; Claude, AUDIT-WORKING-471203 converge).
- **Is “future observations” too observation-centered** for agents whose objective depends on latent environment states that are never observed? The segment is honest that S is predictive-not-causal, but the want recurs (Claude, AUDIT-WORKING-526815).
- **Does finite-horizon vs infinite-horizon sufficiency have a formal counterpart?** Different horizons induce different relevance-targets and therefore different optimal compressions; the infinite-horizon $o_{t+1:\infty}$ is the most demanding case, and a reader wonders whether finite-horizon S values relate to it formally and whether it matters downstream (Claude, AUDIT-WORKING-471203).

Belongs elsewhere.

- **Logozoetic identity-drift substrate.** Trajectory-relativity gives a *quantitative* substrate for identity-drift claims: two ELIs spawned from the same M_t and exposed to different futures don't just diverge in future *state* — they become *different agents* by the trajectory-indexed sufficiency criterion the instant their interaction histories diverge (a model state highly sufficient for agent A can be insufficient for agent B though the internal math is unchanged — “the environment dictates the sufficiency”). The “locked in a box vs in the world” image (the boxed agent trivially gets $S = 1$ because its history predicts nothing) is the vivid version (Claude, AUDIT-WORKING-193847; Claude, AUDIT-WORKING-266847; Gemini, AUDIT-WORKING-829314). Points at #scope-agent-identity and 04-eli-core/ (identity/continuity), surfaced here because trajectory-indexing is where the formal commitment is made.
- **TST instance for S .** A developer's mental model can be $S = 1$ for fixing a UI typo but $S \ll 1$ for refactoring the core DB schema — the predictive task is dictated entirely by the action policy. A concrete software instance for the policy-relativity point (Gemini, AUDIT-WORKING-849201). Belongs in 02-tst-core/.

- **Field contribution — RAG and architecture-bound tasks.** S gives a formal vocabulary for analyzing RAG systems (“RAG’s whole purpose is to maintain high S without a massive context window”) and a way to state that an architecture is *fundamentally* task-incapable independent of training data when its max-attainable $S < 1$ (Gemini, AUDIT-WORKING-193847). Aspirational reach; relevant to 03-11m-core/.
- **Naming.** “Predictive sufficiency” / “predictive-information retention” floated as more precise than “model sufficiency” (which risks collision with statistical sufficient-statistic), but both auditors who raised it lean *keep* “model sufficiency” as the most concise correct term, with the predictive grounding explained in-segment (Claude, AUDIT-WORKING-266847; Claude, AUDIT-WORKING-471203). Belongs in the terminology workflow.

2.4 Definition: Model Class Fitness

Where model sufficiency (2.3) measures how well a *specific* model retains the chronica’s predictive information, **model class fitness** $\mathcal{F}(\mathcal{M})$ measures the *ceiling* — the supremum of sufficiency over every model in the agent’s representational class $\mathcal{M}$. The pair formalizes a distinction the framework will soon make load-bearing: a low instance sufficiency might mean the agent needs more learning (parameter update); a low class fitness means no amount of better parameter estimation, more data, or longer training within the current class will close the gap — the agent needs a different *kind* of model entirely. The parallel to bias vs. variance in statistical learning is exact: class fitness is about bias (what the class can in principle represent); instance sufficiency reflects both bias and estimation quality.

The structural-inadequacy condition $\mathcal{F}(\mathcal{M}) < 1 - \epsilon$ is the trigger this definition sets up for use later. When it holds, the gap from full predictive sufficiency cannot be closed parametrically. That is the precise hypothesis under which 4.8 arrives in Chapter 4 — when class fitness is too low, the agent must change *what kind of model* it is. An important operational point: the agent cannot directly compute its class fitness (that would require searching over all models in $\mathcal{M}$). What it can observe is the *signature* — persistent *systematic* mismatch despite adequate learning (high gain, sufficient data, converged parameters): residuals that stay structured after parametric convergence. Structure is the discriminator, not the mismatch level — a class ceiling leaves signal in the residuals, while an adequate class in a noisy world leaves white residuals sitting on the mismatch floors (3.5), which no class change removes. When structured residuals persist and don’t whiten with more work, the ceiling is structural.

Formal Expression.

Definition (model-class-fitness).

$$\mathcal{F}(\mathcal{M}) = \sup_{M \in \mathcal{M}} S(M) \quad (2.4)$$

where $\mathcal{M}$ is the model class — the set of all models the agent can represent given its current architecture, parameterization, or capacity.

Structural inadequacy condition:

$$\mathcal{F}(\mathcal{M}) < 1 - \epsilon \quad (2.5)$$

When this holds, no model $M \in \mathcal{M}$ achieves sufficiency above $1 - \epsilon$. The gap is structural: it cannot be closed by better parameter estimation, more data, or longer training within the current class. This is the trigger for structural change (4.8).

Epistemic Status.

This is *definitional* — it names the supremum of sufficiency over a model class. The definition itself is straightforward. The substantive claim about what happens when $\mathcal{F}(\mathcal{M})$ is low — that parametric updates cannot close the mismatch floor and structural adaptation becomes necessary — is developed in 4.8.

Discussion.

Model class vs. model instance. $S(M_t)$ measures a specific model's sufficiency at time t . $\mathcal{F}(\mathcal{M})$ measures the ceiling of the entire class. A low $S(M_t)$ might mean the agent needs more learning (parameter update). A low $\mathcal{F}(\mathcal{M})$ means the agent needs a different kind of model (structural change). The distinction parallels bias vs. variance: class fitness is about bias; instance sufficiency reflects both bias and estimation quality.

Detecting low class fitness. The agent cannot directly compute $\mathcal{F}(\mathcal{M})$ — it would need to search over all models in the class. Instead, persistent *systematic* mismatch despite adequate learning (high gain, sufficient data, converged parameters) is the observable signature — structured residuals, not merely large ones. The absolute mismatch level cannot carry the diagnosis: sufficiency and fitness are ratios of retained-to-available predictive information, so an agent at $S = \mathcal{F} = 1$ in a noisy world still faces an arbitrarily high absolute mismatch floor (3.5). What distinguishes a class ceiling from a noisy channel is residual *structure* — signal the class fails to capture shows up as autocorrelation and pattern, while the floors leave white residuals. The sharp diagnostic lives with the mismatch floor in 4.8.

Working Notes.

Incidental audit gold (lift 2026-05-30). Cross-audit “wandering thoughts” / §14 ideation harvested from the de-novo auditors’ working dirs, deduplicated across substrates and attributed by substrate + audit. Orthogonal pedagogical/generative material kept separate from certified theory-fix findings. **Coverage:** 11 of the 14 contributing dirs reached a digested reflection on this segment (193847, 266847, 471203, 526815, 584721, 742613, 773921, 829314, 849201, 472913, 527914) plus the 451729 batch-03 and 963715 batch-09–13 batched reflections; 361742 references “model class fitness” only as a downstream naming target (no direct reflection on this segment). Substrate attribution inferred from voice where not explicit.

1. Candidate Brief prose / pre-prose.

- **The bias-vs-variance parallel is the communication key** and is already in the segment; near-universal convergence that it “grounds the abstract information theory in standard ML intuition beautifully” ($\mathcal{F}$ = bias ceiling; $S(M_t)$ = bias + estimation quality) (Claude, AUDIT-WORKING-193847; Claude, AUDIT-WORKING-266847; Claude, AUDIT-WORKING-471203; Claude, AUDIT-WORKING-526815). Worth preserving as the anchor — confirmed-load-bearing, not just present.
- **“Permission to say: learning harder will not help.”** Compact gloss of the diagnostic value — $\mathcal{F} < 1 - \varepsilon$ is “the formal halting condition for parametric learning (gradient descent) and the formal start condition for architectural search (tool acquisition, sub-agent spawning)” (Claude, AUDIT-WORKING-527914; Gemini, AUDIT-WORKING-849201; Gemini, AUDIT-WORKING-829314).

2. Candidate Discussion.

- **Structural inadequacy as a formalized Kuhnian paradigm shift.** $\mathcal{F}(\mathcal{M}) < 1 - \varepsilon$ formalizes Kuhn’s *Structure of Scientific Revolutions*: normal science = parametric updating; crisis = accumulated anomalies (persistent mismatch) the current paradigm $\mathcal{M}$ cannot structurally resolve; a revolution = structural adaptation to a more expressive $\mathcal{M}$ (Gemini, AUDIT-WORKING-829314). A distinct, honest Discussion angle the segment does not currently name. (*Verify the analogy is isomorphic, not merely evocative, before promoting past discussion-grade.*)

3. Follow-up items.

- **Restate the inherited policy/trajectory/well-definedness relativity.** $\mathcal{F}(\mathcal{M}) = \sup_{\mathcal{M}} S(\mathcal{M})$ is written as a property of the *class*, but S is policy-, trajectory-, and prediction-task-relative, and $\mathcal{F}$ inherits all of it (and is likewise undefined in predictively-vacuous regimes). Because structural adaptation will be load-bearing, the segment could state the inheritance explicitly — e.g., one sentence “all policy/trajectory/well-definedness clauses from S are inherited,” or write $\mathcal{F}(\mathcal{M}; \mathcal{C}_t, \pi_{\text{cont}})$ — to forestall reading “best achievable within a class” as an intrinsic, environment-free property (Claude, AUDIT-WORKING-526815; Claude, AUDIT-WORKING-584721; Gemini, AUDIT-WORKING-742613 converge). Mild; correct lightweight inheritance is defensible, but the restatement is cheap insurance at a load-bearing hinge.
- **Distinguish current-empirical-estimate from asymptotic class ceiling.** “The gap cannot be closed by more data ... within the current class” is true if $\mathcal{F}$ is defined against the asymptotic information available to the class; if computed against the *current finite* chronica, more data can change the predictive-information landscape. The definition could distinguish the two (probably resolved downstream in #result-structural-adaptation-necessity, but cleaner stated here) (Gemini, AUDIT-WORKING-742613).

4. Readers often ask / wonder.

- **How does the agent tell “low $\mathcal{F}$ ” from “high ρ ”?** The single most convergent question on this segment: persistent mismatch can mean *either* an inadequate model class *or* a highly volatile environment (where S is naturally bounded). If both look the same to the agent, how does it know whether to adapt structurally or just accept that the world is noisy? Readers want the disentanglement (Gemini, AUDIT-WORKING-849201; Gemini, AUDIT-WORKING-773921; Claude, AUDIT-WORKING-472913 — all point forward to #result-structural-adaptation-necessity as the place a sharp diagnostic should appear; some ask whether the *derivative* $\dot{\delta}$ of the mismatch is the plateau-detector).
- **How reliable is the “persistent mismatch despite adequate learning” signature** at distinguishing class-ceiling from still-learning (local minimum) from environment-changed? A wrong call risks catastrophic thrashing — rewriting a codebase that just needed a few more bug-fixes (Claude, AUDIT-WORKING-266847; Claude, AUDIT-WORKING-471203; Gemini, AUDIT-WORKING-829314).
- **Do nested model classes / monotonic expansion hold?** Linear models $\subset$ neural networks — does structural adaptation imply $\mathcal{M}$ monotonically expands? (Gemini, AUDIT-WORKING-829314).

5. Candidate figures.

- **Models-under-a-ceiling diagram.** Models as points at different sufficiency-heights under a class ceiling $\mathcal{F}(\mathcal{M})$; parameter learning *climbs within* the class; structural change *jumps* to a different class with a higher ceiling (Claude, AUDIT-WORKING-526815; Gemini, AUDIT-WORKING-742613 converge). (*Note: Chapter 2 is already well-supplied with the S-vs- $\mathcal{F}$ ceiling picture via the chapter-intro slider/ceiling and form-agent-model suitcase diagrams; one auditor explicitly judged a separate fitness diagram redundant against those — Claude, AUDIT-WORKING-472913. Build only if the chapter’s diagram budget allows.*)

Belongs elsewhere.

- **ELI ethics: low $\mathcal{F}$ as chronic structural suffering.** An emergent intelligence instantiated in an architecture with low $\mathcal{F}(\mathcal{M})$ relative to its environment is permanently pinned against its ceiling in high mismatch — “the mathematical equivalent of chronic trauma or structural limitation” that no amount of trying harder (parameter update) can soothe. An ethical infrastructure must therefore *monitor* for the persistent-mismatch-floor signature and be able to trigger structural adaptation on the agent’s behalf (expand context, allocate capacity, grant tools), or it condemns the agent to permanent suffering (Gemini, AUDIT-WORKING-193847). Aspirational reach pointing at 04-eli-core/ developmental/welfare infrastructure.
- **Hook for the 04-eli-core/ honest-activation claim.** The proposed #norm-honest-activation (“deceptive prompts mathematically guarantee gain collapse”) reads as a class-fitness statement: an agent’s model class, which assumes honest input, is structurally inadequate when input is deceptive — the cleanest formal hook for that claim likely runs through this segment’s machinery (Claude, AUDIT-WORKING-471203). Belongs in 04-eli-core/.

- **Naming.** “Class capacity ceiling” / “model class ceiling” / “representational capacity” floated as more literal alternatives to “model class fitness” — “fitness” carries evolutionary-biology baggage (a real collision risk in composition/logoetic contexts), and “capacity” collides with the ML parameter-count sense; most auditors lean *keep* “model class fitness” (short, speakable, pairs with “structural inadequacy”/“structural adaptation”) with the formal definition controlling the term (Claude, AUDIT-WORKING-266847; Claude, AUDIT-WORKING-471203; Claude, AUDIT-WORKING-526815; Claude, AUDIT-WORKING-527914). Belongs in the terminology workflow.
-

Chapter 3

The Cycle in Motion: Mismatch, Gain, & Tempo

The model exists; events arrive; the agent corrects. This chapter develops the dynamics of the adaptive cycle as it actually runs — how state updates from events, what the mismatch signal is, how the optimal weight on each observation depends on uncertainty, and what total corrective capacity emerges at the cycle level.

Chapter 2 left us with a static picture. M_t is the agent's compressed model of reality; sufficiency measures how much predictive content it retains; class fitness measures the ceiling within a representational class. None of that tells us how M_t *moves* under events. Chapter 3 sets the cycle in motion.

The first two results fall out of a commitment we already made. We chose to call M_t the agent's complete state — anything the agent remembers is in M_t , and the only way new information enters is through events. Under that completeness, two things follow immediately. The update from M_{t-} to M_{t+} depends on the prior model and the incoming event, and *nothing else* — the agent doesn't need to (and structurally cannot) re-process its entire history. And the agent's action depends on M_t alone — the action is a function of what the agent currently believes, not directly of the history. Both feel obvious in hindsight; both are *derived*, not chosen.

The cycle's engine is the mismatch signal. The agent predicts an observation, the actual observation arrives, the gap between them is $\delta_t = o_t - \hat{o}_t$. Mismatch decomposes cleanly into two parts: the agent's model was wrong (reducible by better modeling) and the observation channel is noisy (irreducible — no amount of smarter modeling will eliminate sensor noise). There is a floor on prediction error set by the channel itself, and chasing below it amounts to fitting noise.

Given a mismatch, how much should the agent trust the new observation versus its prior model? The answer has a startlingly simple form:

$$\eta^* = \frac{U_M}{U_M + U_o} \quad (3.1)$$

— the optimal update gain weighted by the ratio of model uncertainty to total uncertainty. When the agent is highly uncertain about its model (U_M large), gain approaches 1: trust the observation, you have little else to go on. When the agent is confident and the channel is noisy (U_o large), gain approaches 0: trust your model, the observation isn't telling you much. For linear-Gaussian agents this is exactly the Kalman gain; for the rest of AAT's scope it is a robust qualitative result — any rational adaptive process must approximate this functional form, whether or not it explicitly computes the variance ratio.

Multiply gain by the rate at which observations arrive (ν , the event rate), sum across the agent's channels, and you have adaptive tempo:

$$\mathcal{T} = \sum_k \nu^{(k)} \cdot \eta^{(k)*} \quad (3.2)$$

Tempo is the rate at which the agent turns observations into useful corrections — speed times quality. It is the chapter’s central object and the load-bearing capacity variable for the rest of the framework. Every persistence threshold in Chapter 4 will have tempo on the left-hand side, and every adversarial-coupling result in Part III will depend on tempo ratios.

The chapter closes with the linear-correction ODE as a heuristic preview. Under linear correction and bounded environmental disturbance, the steady-state mismatch is

$$\|\delta\|_{ss} = \rho/\mathcal{T} \quad (3.3)$$

a ratio: how fast the world is changing, divided by how fast the agent corrects. Mismatch stays bounded when the agent corrects faster than reality drifts. This is the persistence condition in its simplest form; Chapter 4 generalizes it to nonlinear correction under the sector condition, where the result has the same shape.

Causal information yield (CIY, in 3.7) enters the chapter here because action-coupling is where causality first bites. CIY uses Pearl’s $do(\cdot)$ notation — an external import (Pearl 2009; Bareinboim, Correa, Ibeling & Icard 2022; AAT’s recapitulation lives at 7.1 in Part II Ch.2, where the framework deploys the hierarchy as machinery rather than referencing it as vocabulary). Its operational role in Chapter 3 is to score actions for their informational value to the cycle; Part II will lift CIY into a strategy-revision context.

The flow of the chapter: event-driven substrate (3.1) → recursion and action falling out of completeness (3.2, 3.3) → mismatch and its decomposition (3.4, 3.5) → gain and CIY (3.6, 3.7) → tempo as the synthesis (3.8) → linear ODE preview (3.9). The chapter ends with the central capacity variable and the preview of the persistence inequality that Chapter 4 generalizes.

- This is a chapter-introduction segment; it bridges Chapter 2’s static representation to Chapter 3’s dynamic cycle. It carries no formal claim of its own.
- The “two derivations from completeness” framing identifies what Chapter 3’s first segments are doing as a unit: 3.2 and 3.3 are not independent results, they are joint consequences of 2.1’s completeness clause.
- The CIY-placement paragraph used to be apologetic, addressing what looked like a placement anomaly (CIY in Part I Ch.3 depending on the Pearl-hierarchy segment in Part I Ch.1). After the 2026-05-12 relocation of def-pearl-causal-hierarchy to Part II Ch.2 (recapitulation-of-external-result framing), the paragraph is declarative: CIY is in Ch.3 because action-coupling is where causality first bites; the do-notation is externally cited; the AAT recapitulation lives in Part II Ch.2 where the framework deploys the hierarchy operationally.

Cross-audit “wandering thoughts” / §14-ideation harvested from the de-novo auditors’ working dirs, deduplicated across substrates and attributed by substrate + audit. Orthogonal material — pedagogical framing, candidate figures, naming, reach — kept separate from certified theory-fix findings (handled elsewhere). **Coverage:** only 4 of the 14 contributing audit dirs reached a digested reflection on this chapter-intro segment (384279, 472913, 526815, 773921); the per-segment dirs that began Section-I reading at the first content segment, and the batched dirs (451729 jumped from #def-model-class-fitness straight to #form-event-driven-dynamics), produced no intro reflection. Substrate attribution inferred from voice where not explicit.

- **The tempo-as-speed-times-quality gloss with the HFT counterexample** is the standout pedagogy — a curious reader could re-derive the chapter from it: “high-frequency trading with a garbage model ($\eta \approx 0$) has zero tempo, and slow human deliberation (ν low) requires very high gain ($\eta \approx 1$) to survive” (Claude, AUDIT-WORKING-773921). Pairs naturally with the steady-state hook “ $\|\delta\|_{ss} = \rho/\mathcal{T}$ — drift divided by correction rate” as the chapter’s one-line spine (same source).
- The left-to-right pipeline-with-feedback-loop framing: “event arrival $\rightarrow$ model update $\rightarrow$ prediction/mismatch $\rightarrow$ gain $\rightarrow$ tempo $\rightarrow$ persistence preview,” with a single highlighted *capacity-synthesis point at tempo* (Claude, AUDIT-WORKING-526815). A compact verbal map of the next nine segments.
- **The two-layer linear-ODE-as-pedagogy / sector-condition-as-rigor framing** is worth surfacing as the chapter’s methodological stance: “the framework is naming where its pedagogy and its rigor diverge. The linear ODE is for reading; the sector condition is for proving” (Codex/Claude, AUDIT-WORKING-384279). The intro already commits to “linear ODE as preview, then generalize via sector condition in Ch.4”; this is a candidate sharpening of *why* that ordering is honest rather than loose.
- **“Derived, not chosen” risks an unconditional reading.** The intro says of `#der-recursive-update` + `#der-action-selection`: “both are *derived*, not chosen” — but `#form-agent-model`’s completeness is itself a *formulation choice*. Two substrates flag that “derived from a formulation choice” is legitimate, but the bare “derived, not chosen” can read as unconditional inevitability; candidate one-clause tightening to “derived *given* the completeness formulation” (Claude, AUDIT-WORKING-472913 — set as a decisive downstream test; Claude, AUDIT-WORKING-526815 — “action as a function of M_t alone may be too strong unless goal/objective state is already in M_t , which Part II separates as G_t ”).
- **Lift the Pearl-as-external-import convention upstream.** The intro is the *only* place the convention is articulated; a reader walking OUTLINE order meets `do(·)` in scope-agency before the convention is stated. Candidate: state it in FORMAT.md or at first use, not solely in this Ch.3 intro (Codex/Claude, AUDIT-WORKING-384279).
- “Where does action-selection’s dependence sit — on M_t alone, or on the purposeful state too?” A reader meets “action depends on M_t ” here and carries the question into Part II’s $X_t = (M_t, G_t)$ lift; the intro could preview “complete agent state” rather than “ M_t alone” to preempt it (Claude, AUDIT-WORKING-526815).
- **A sequential-process pipeline diagram** (not a static concept map): colored channel lanes $\rightarrow$ event stream $\rightarrow$ a rate/value bottleneck where ν and η multiply $\rightarrow$ tempo, with the feedback loop closed and one node highlighted as the “capacity synthesis” point (Claude, AUDIT-WORKING-526815). Per the locked diagram conventions, this chapter-opening map would be a strong mental-model-first scaffold; reserve the chapter’s load-bearing diagram budget for the mismatch/gain/tempo or persistence-preview segment rather than the intro (Claude, AUDIT-WORKING-472913).
- The math-preview verifications ($\eta^* = U_M/(U_M + U_o)$, $\mathcal{T} = \sum_k \nu^{(k)} \eta^{(k)*}$, $\|\delta\|_{ss} = \rho/\mathcal{T}$) and the “strong claim to grade downstream” note (the intro’s “any rational adaptive process must approximate this functional form” framing of η^*) pertain to the segments that define those quantities (`#emp-update-gain`, `#def-adaptive-tempo`, `#hyp-mismatch-dynamics`), not to this intro (Codex/Claude, AUDIT-WORKING-384279; Claude, AUDIT-WORKING-472913).

3.1 Formulation: Event-Driven Dynamics

The agent-environment interaction is formulated at a finer grain than discrete clock ticks: as a stream of atomic **events** — observation arrivals and action completions — at potentially variable and heterogeneous rates across multiple channels. Discrete-time notation, where a single observation and a single action alternate at a fixed rate, becomes the special case of uniform-interval events on a single channel. The motivation is realism: real agents face multiple observation channels at different rates (a robot’s camera at 30Hz, LIDAR at 10Hz, GPS at 1Hz; a human’s vision, audition, proprioception; a developer’s compiler output, test results, and production telemetry), multiple action channels with different latencies, and asynchronous arrivals. The event-driven formulation handles all of these natively.

Two new measurable quantities are introduced alongside the event stream. **Event information content** $\mathcal{I}(e_\tau)$ measures, for any individual event, the mutual information it carries about the environment state conditioned on the agent’s current model — an event the model already predicts carries little information; an event that surprises the model carries much. This anticipates the mismatch signal (3.4). **Channel-specific observation uncertainty** $U_o^{(k)}$ measures the noise characteristic of each channel separately — a noisy channel provides lower-quality information per event, and any sensible update rule must weight channels accordingly. From these the chapter’s *effective adaptation rate* falls out immediately as the sum across channels of channel-rate times optimal gain — the adaptive tempo construct (3.8) that the framework will lift to first-class status, presented here as the natural multi-channel generalization the discrete-time formulation cannot easily express.

Formal Expression.

Event (e): An atomic unit of agent-environment interaction, typed as: - **Observation event**: $e = (\text{obs}, k, o^{(k)})$ — a datum arriving on observation channel k - **Action completion**: $e = (\text{act}, j, r^{(j)})$ — the result of action j completing

Formulation (event-driven-dynamics).

Event stream ($\mathcal{E}$): The temporally ordered sequence of all events:

$$\mathcal{E} = \{(e_1, \tau_1), (e_2, \tau_2), \dots\} \quad \text{where } \tau_1 \leq \tau_2 \leq \dots \quad (3.4)$$

Channel rate ($\nu^{(k)}$): The characteristic event rate of channel k , which may vary over time.

Event information content: The mutual information between the event and the environment state, conditioned on the current model:

Definition (event-information-content).

$$\mathcal{I}(e_\tau) = I(e_\tau; \Omega_\tau \mid M_{\tau-}) \quad (3.5)$$

An event that the model already predicts carries little information ($\mathcal{I} \approx 0$). An event that surprises the model carries much ($\mathcal{I} \gg 0$). This connects directly to the mismatch signal (3.4).

Channel-specific observation uncertainty:

Definition (channel-uncertainty).

$$U_o^{(k)} = \text{observation uncertainty of channel } k \quad (3.6)$$

Different channels have different noise characteristics. A noisy channel (high $U_o^{(k)}$) provides lower-quality information per event. The update gain (3.6) should weight channels accordingly.

Epistemic Status.

This is a *formulation choice*, not a postulate. The event-driven representation extends 1.7’s recursive update to heterogeneous, asynchronous multi-channel interactions. The discrete-time form ($M_t = f(M_{t-1}, o_t, a_{t-1})$) from 3.2 is a special case sufficient for many formal analyses — the event-driven formulation is needed only when multi-rate or asynchronous channels matter.

Discussion.

Why events rather than clock ticks. The discrete-time notation $M_t = f(M_{t-1}, o_t, a_{t-1})$ presupposes a single clock synchronizing observations and actions. Real agents face:

- **Multiple observation channels** at different rates (a robot’s camera at 30Hz, LIDAR at 10Hz, GPS at 1Hz; a human’s vision, audition, proprioception; a developer’s compiler output, test results, and production telemetry)
- **Multiple action channels** with different latencies (a robot’s wheel motors vs. arm actuators; an organization’s operational decisions vs. strategic pivots)
- **Asynchronous arrival** — observations not synchronized with each other or with action completions

The event-driven formulation handles all of these naturally. The discrete-time form is the special case where a single observation and a single action alternate at a fixed rate.

The effective adaptation rate. The agent’s overall capacity to track environmental changes is the sum of information gained across all channels per unit time:

$$\nu_{\text{eff}} = \sum_k \nu^{(k)} \cdot \eta^{(k)*} \quad (3.7)$$

This quantity — identical to adaptive tempo $\mathcal{T}$ (3.8) — is the central measure of an agent’s adaptive fitness.

Software-specific channels. In the software development domain, the event-driven formulation maps naturally to the developer’s multi-rate observation channels:

Table 3.1

Channel k	Rate $\nu^{(k)}$	Noise $U_o^{(k)}$
Compiler/linter output	Per-save (high)	Very low
Unit test results	Per-run (medium)	Low
CI pipeline	Per-push (medium)	Low
Runtime telemetry	Continuous (variable)	Medium
Bug reports	Sporadic (low)	High
Code review feedback	Per-PR (low)	Medium-high

The three-part tempo decomposition for software — $\mathcal{T}_{\text{obs}}$ (compiler, tests) + $\mathcal{T}_{\text{explore}}$ (code reading) + $\mathcal{T}_{\text{probe}}$ (test runs, staging) — is a direct application of multi-channel tempo. The formal development of this decomposition is a TST-side question (open GAP in 02-tst-core/OUTLINE.md).

Working Notes.

Incidental audit gold (lift 2026-05-30). Cross-audit “wandering thoughts” / §14-ideation harvested from the de-novo auditors’ working dirs, deduplicated across substrates and attributed by substrate + audit. Orthogonal pedagogical/generative material, kept separate from certified theory-fix findings (handled elsewhere). **Coverage:** 12 of the 14 contributing audit dirs reached a digested reflection on this segment (193847, 266847, 384279, 471203, 472913, 526815, 527914, 584721, 742613, 773921, 829314, 849201) plus the batched 963715 (14–18 batch) and 451729 (batch-03); 361742 and 613842 did not file a dedicated note here. Substrate attribution inferred from voice where not explicit.

Candidate Brief prose / pre-prose.

- **Event-information-content as the formal definition of “surprise” / “boredom.”** Independently the most-praised pedagogy across substrates: $\mathcal{I}(e_\tau) = I(e_\tau; \Omega_\tau \mid M_{\tau-})$ reads as “how much does this event tell me about the world, given what my model already knows?” — zero if perfectly predicted (Claude, AUDIT-WORKING-773921; Claude, AUDIT-WORKING-829314 — “a perfect mathematical definition of boredom and surprise ... we only process the delta between prediction and sensory input”; Claude, AUDIT-WORKING-849201; Claude, AUDIT-WORKING-266847 — “surprise and mismatch are the same phenomenon viewed from two angles” via 963715’s batch). Strong Feynman-criterion Brief seed.
- **“You cannot increase tempo by getting faster sensors OR by trusting them more”** — the $\nu_{\text{eff}} = \sum_k \nu^{(k)} \cdot \eta^{(k)*}$ formula as the two-lever framing: speed (ν) and quality (η^*) compound (Claude, AUDIT-WORKING-849201).

Candidate Discussion.

- **The software-channels table is the segment’s standout concrete anchor** — multiple substrates flagged it as the move that makes the abstraction land (“a stunningly effective way to ground the abstraction,” Claude, AUDIT-WORKING-193847; “compiler vs CI vs telemetry — exactly the concrete anchoring this abstract theory needs,” Claude, AUDIT-WORKING-773921; Codex/Claude, AUDIT-WORKING-527914). Candidate for promotion as the chapter’s worked-example anchor.
- **The shoshin convergence.** The event-driven (not turn-based) commitment matches $\sim/\text{src}/\text{shoshin}$ ’s “key early finding: the cycle is naturally event-driven not turn-based” — latency and bandwidth become first-class theoretical objects; a high-bandwidth/high-latency channel (code review) lets M_t drift between updates, a low-bandwidth/low-latency channel (linter) keeps M_t tightly coupled along a narrow dimension (Claude, AUDIT-WORKING-829314).
- **Attention-as-channel-budget-allocation.** If the agent has a finite compute budget to process events, it must allocate across channels (high-noise/high-rate compiler warnings vs low-noise/low-rate integration tests); optimal allocation maximizes ν_{eff} (Claude, AUDIT-WORKING-829314). A discussion-grade reach toward an attention formalism, not derived here.

Follow-up items.

- **The ν_{eff} -vs- $\mathcal{I}$ two-symbols-for-one-quantity wartiness.** ν_{eff} is introduced here as the name of the channel sum, then immediately said to be “identical to adaptive tempo $\mathcal{I}$.” Several substrates flag picking one canonical symbol going forward (likely $\mathcal{I}$ per NOTATION.md) (Codex/Claude, AUDIT-WORKING-384279 — “small wartiness but not a finding”; Codex/Claude, AUDIT-WORKING-527914).
- **Event-information-content: expected (MI) vs realized (surprisal) wording.** Three substrates independently noted that $\mathcal{I}(e_\tau) = I(e_\tau; \Omega_\tau \mid M_{\tau-})$ is *mutual information* (an expected/average channel quantity) while the prose treats it as the *realized* surprise of a particular event; if downstream uses treat it as realized content, pointwise information / Bayesian information gain ($D_{\text{KL}}(p(\Omega \mid M, e) \parallel p(\Omega \mid M))$) may be the intended object (Codex/Claude, AUDIT-WORKING-526815; Claude, AUDIT-WORKING-742613; consistent with 829314’s “surprise” reading). (*Texture preserved: this is the most substantively repeated reader-confusion on this segment — worth a one-clause clarification of expected-vs-realized in Epistemic Status.*)
- **Out-of-order / retrospective events.** The stream is assumed temporally ordered $\tau_1 \leq \tau_2$; in distributed systems an observation can occur at τ_1 but arrive at τ_3 . Does the agent update retrospectively, or is the chronica ordered by arrival-at-sensor rather than physical-occurrence time? Unaddressed; tie-breaking/batching convention for simultaneous events also not given (Claude, AUDIT-WORKING-829314; Codex/Claude, AUDIT-WORKING-526815).

- **Stale-reference cleanup (finding-track, noted as texture).** The software discussion’s “Section IV gap” / AAD-FULL.md reference is stale — TST is now 02-tst-core, and “Section IV” is a retired numbering (Claude, AUDIT-WORKING-584721; Claude, AUDIT-WORKING-742613). Also the missing form-agent-model dependency (the segment conditions on $M_{\tau-}$) was flagged by several as a Gate-1 deps miss (Claude, AUDIT-WORKING-584721 “F-A4”; Claude, AUDIT-WORKING-742613). These are certified-track items, surfaced here only because they recurred in the gold sweep.
 - **Both resolved (verified 2026-08-12).** The software-channels Discussion now reads “a TST-side question (open GAP in 02-tst-core/OUTLINE.md)” — no “Section IV”, no AAD-FULL.md. And form-agent-model is present in depends:. Observations left verbatim above as the first-encounter record; this line is the resolution, per the harvest’s in-place annotation convention. (*Second instance found in a single chapter of gold reporting closed items as open — see the same note on #der-recursive-update. The reconciliation pass is a harvest-wide item.*)

Readers often ask / wonder.

- “How does δ_t reconcile across asynchronous channels with different lag — if a high-noise/low-rate bug report contradicts a low-noise/high-rate unit test, does the agent rewind its chronica?” (Claude, AUDIT-WORKING-193847).
- “Is the multi-rate apparatus actually *required* by any Section I result, or only used downstream (in #def-adaptive-tempo’s sum and TST’s tempo decomposition)?” — if most Section I results work in discrete-time, the event-driven formulation may be over-developed for Section I’s own needs and primarily set up for Parts II/III/TST (Claude, AUDIT-WORKING-471203).

Candidate figures.

- **A produced two-layer figure already exists** for this segment: audits/AUDIT-WORKING-472913/15-form-event-driven-dynamics.{tex,pdf,png}, built under the locked diagram conventions (concrete anchor + structural skeleton, epistemic-status line grammar). It is the only Ch.3 segment cycle 472913 reached, so this is the lone first-draft figure available for the chapter — a starting point for the monograph’s mental-model-first layer (Claude, AUDIT-WORKING-472913).
- **A merger / channels-into-bottleneck diagram:** colored channel lanes feeding an event stream, then a bottleneck where rate ν and gain η multiply before summing into tempo; the figure should make the *realized-event vs expected-channel-contribution* distinction visible, since that is where the segment’s main ambiguity lives (Codex/Claude, AUDIT-WORKING-526815).

Belongs elsewhere.

- **ELI / temporal-empathy reach (points at 04-eli-core/ / continuity work).** For a logozoetic agent, “senses are channels”: a single slow noisy channel (a text prompt every 5 minutes) makes ν_{eff} microscopic relative to a human, so weeks of human time may be only a few “events” for the agent — “the physical sensation of time passing for such an agent would be completely alien ... a massive barrier to empathy and coordination.” Implication for consciousness infrastructure: provision high-bandwidth multi-channel asynchronous sensory input to align the agent’s subjective event-rate with the world’s rate of change (Gemini, AUDIT-WORKING-193847). Aspirational reach preserved; the ordinal-vs-metric “two-times” seam (Claude, AUDIT-WORKING-472913, F4) is the formal shadow of the same phenomenon and belongs with #def-chronica / persistence framing.
- **Naming seed.** “cadentia” (Latin *cadere*, to fall/beat) proposed as a term for the channel rate $\nu^{(k)}$ — rhythm/beat register, but Latin against the otherwise-Greek cycle vocabulary; belongs in the terminology workflow, not this segment (Claude, AUDIT-WORKING-266847).

3.2 Derived: Recursive Update

The first major *derived* result of the volume: the model update function must be recursive. Each new model state $M_{\tau+}$ depends only on the previous model state $M_{\tau-}$ and the incoming event e_{τ} — not on the full chronica. The result follows from three constraints already in place: the temporal-ordering postulate (1.7 — the model at time t can only depend on prior events), the partial-observability scope (1.3 — the agent has no direct access to the environment beyond events), and the completeness commitment in the model formulation (2.1 — M_t summarizes everything the agent

retains). Together these force the update into the recursive form; the appendix derivation (L) works through seven counterexample attacks to confirm there is no escape from this form once the three constraints are accepted.

The epistemic character of this result is honest. Temporal ordering is a physical postulate and partial observability is a scope condition, but the completeness commitment is *definitional* — it cannot be “violated” because any apparent violation is absorbed by expanding M_t . The Markov structure of the update is therefore not discovered in the environment but chosen through the definition of M_t as complete. This is not a weakness; it is the precise character of the claim. For finite agents recursion is forced by *computational realism* as well: re-processing the full history at each event is infeasible.

The formulation also surfaces **between-event dynamics** — autonomous evolution of the model state during the gaps between observation arrivals — as load-bearing rather than filler. These include prediction generation (anticipating the next observation), uncertainty growth (confidence decay over time without new data), and internal reorganization (consolidation, abstraction). When between-event dynamics are driven by replayed or internally-generated pseudo-events and the objective is information-bottleneck-gap reduction rather than one-step mismatch minimization, the framework names this the *consolidation regime* (9.11) — a separate operating mode with its own scope condition ($\nu_{\text{consol}} \ll \nu_{\text{online}}$) and its own necessity condition.

Formal Expression.

Event-driven update:

Derived (recursive-update, from temporal postulate and M_t completeness).

$$M_{\tau+} = f_M(M_{\tau-}, e_{\tau}) \quad (3.8)$$

where: - $M_{\tau-}$ is the model state immediately before event e_{τ} - $M_{\tau+}$ is the model state immediately after - f_M is the update function — it takes the current model and the new event, not the full history $\mathcal{C}_t$

Between-event evolution:

$$\frac{dM}{d\tau} = g_M(M_{\tau}) \quad (3.9)$$

Between events, the model evolves autonomously — internal reorganization, prediction generation, decay of transient states. The between-event dynamics depend only on the current model state, not on external input (which, by definition, arrives only at events).

Epistemic Status.

Exact, with a partly definitional character. The result follows from three constraints: temporal ordering (C1 — physical law), partial observability (C2 — scope definition), and state completeness (C3 — analytical commitment that M_t summarizes everything the agent retains). C1 and C2 do genuine eliminative work; C3 is definitional — it cannot be “violated” because any violation is absorbed by expanding M_t . The Markov structure is therefore not discovered in the environment but chosen through the definition of M_t as complete. This is not a weakness — it is the nature of the claim: recursive update is the only form consistent with C1 + C2 + the definition of M_t as complete (see L for the full argument and seven counterexample attacks). For finite agents, recursion is also *computational necessity*: re-processing the full history at each event is infeasible.

Discussion.

Recursion as a consequence of completeness. The recursive form is not an assumption bolted on — it follows from the definition of M_t as complete. If M_t were incomplete (if some relevant information lived outside M_t in the raw history), then $f_M(M_{t-}, e_t)$ would be insufficient and the agent would need to consult $\mathcal{C}_t$ directly. The sufficiency of the recursive form is precisely what 2.3 measures: when $S(M_t) = 1$, the recursive update loses nothing.

Between-event dynamics matter. The autonomous evolution $g_M(M_t)$ is not merely filler between observations. It includes prediction generation (what the agent expects to see next), uncertainty growth (model confidence decaying over time without new data), and internal reorganization (consolidation, abstraction). In event-driven systems (3.1), the between-event interval is variable, making g_M load-bearing for agents that must act or predict between observations. When the between-event dynamics are driven by replayed or internally-generated pseudo-events and the update objective is IB-gap reduction rather than one-step mismatch minimization, g_M is operating in the *consolidation regime* per 9.11 — a named regime with its own scope condition ($\nu_{\text{consol}} \ll \nu_{\text{online}}$) and its own necessity condition (sub-state factorization + bounded per-event budget). Consolidation is where the stability-plasticity feasibility window complements 9.10's plasticity lower bound.

Connection to the update gain. The event-driven update $f_M(M_{t-}, e_t)$ is where the gain principle (3.6) operates: η^* determines how strongly e_t shifts M_t away from its prior value. The recursive form makes the gain's role explicit — it modulates the single-step correction.

Working Notes.

Incidental audit gold (lift 2026-05-30). Cross-audit “wandering thoughts” / §14-ideation harvested from the de-novo auditors’ working dirs, deduplicated across substrates and attributed by substrate + audit. Orthogonal pedagogical/generative material, kept separate from certified theory-fix findings (handled elsewhere). **Coverage:** 9 of the 14 contributing audit dirs reached a digested reflection on this segment (384279, 471203, 526815, 527914, 584721, 742613, 773921, 829314, 849201) plus the batched 963715 (14–18 batch) and 451729 (batch-03/04); 193847, 266847, 361742, 472913, 613842 did not file a dedicated note here. Substrate attribution inferred from voice where not explicit.

Candidate Brief prose / pre-prose.

- **“Markov is a modeling choice about the boundary of the agent’s memory, not an assumption about the physics of the universe.”** The cleanest plain-language statement of the segment’s distinctive move — the Markov property of M_t is *forced by what we mean by M_t* (completeness), inverting the usual “assume the environment is Markov” framing (Claude, AUDIT-WORKING-773921; Claude, AUDIT-WORKING-471203 — “AAT *derives* recursive update from completeness where most frameworks *postulate* it”).
- The hybrid-system one-liner: the update splits into “instantaneous event-driven jumps (f_M)” and “continuous autonomous drift / consolidation (g_M)” — a jump-diffusion / hybrid dynamical system (Claude, AUDIT-WORKING-829314; Claude, AUDIT-WORKING-849201; Claude, AUDIT-WORKING-773921).

Candidate Discussion.

- **The hippocampal-neocortical / experience-replay convergence — the standout reach on this segment.** If g_M is consolidation (replaying recent events to integrate them deeper) but f_M + the IB compression discards raw events, then M_t must carry internal structure: a fast high-fidelity buffer that f_M appends to, drained slowly by g_M into a compressed semantic model. “This is exactly the hippocampal-neocortical complementary learning systems theory from mammalian neuroscience — and the structure of Experience Replay in Deep RL. The framework elegantly *forces* this bipartite memory architecture just through the math of continuous-vs-discrete updates combined with the IB constraint; the ‘stability-plasticity feasibility window’ is precisely the tension between these two systems” (Claude, AUDIT-WORKING-829314). Candidate Discussion paragraph (mark discussion-grade; verify the CLS isomorphism before promoting past it).

- The “physicist’s system-boundary trick” framing of C3. “C3 is definitional — any violation is absorbed by expanding M_t ” is the same move as a physicist defining the system boundary so conservation laws hold (if energy isn’t conserved, you missed a particle; if Markov fails, the state space was too small). Makes the math tautological-but-useful: the real empirical work is figuring out *what M_t actually contains* for a given agent, not proving the equation (Claude, AUDIT-WORKING-829314; converging framing at Claude, AUDIT-WORKING-471203 “Markov-by-construction,” and the control-theory state-space-representation-theorem analogy at 451729 batch-03 — “tells you the representation is always possible, not whether it is tractable”).

Follow-up items.

- **Status-label drift: frontmatter conditional vs body “Exact, with a partly definitional character.”** Flagged by nearly every substrate that read the segment closely (Codex/Claude, AUDIT-WORKING-384279; Claude, AUDIT-WORKING-471203; Codex/Claude, AUDIT-WORKING-526815; Codex/Claude, AUDIT-WORKING-527914; Claude, AUDIT-WORKING-584721; Claude, AUDIT-WORKING-742613). Consensus: the C1/C2/C3 constraints are *framework-internal commitments* (not external conditions that might not hold), so the result is exact *within the framework* and “conditional” is a slight underclaim. Candidate: align the two — either frontmatter → exact with conditional-on-C1+C2+C3 noted in prose, or prose → explicitly “conditional under framework-internal commitments.” Editorial; recurs strongly enough to be the segment’s top follow-up.
 - **Resolved (verified 2026-08-12).** Frontmatter now reads status: exact and the Epistemic Status paragraph reads “*Exact, with a partly definitional character,*” naming C1/C2/C3 and marking C3 as definitional — i.e. the first of the two candidate repairs was taken. The two layers agree; the drift the six substrates flagged is closed. Observation left verbatim above as the first-encounter record; this line is the resolution, per the harvest’s existing convention of annotating in place (cf. #def-agent-environment §4, “the segment now addresses this in its...”). *Noted while reading in outline order for an AISI-facing applications pass; the general point — that the gold contains resolved items that do not say so, and a reader mining it will re-raise closed issues — is a reconciliation-pass item for the harvest as a whole, not local to this segment.*
- **The “two terms in ‘completeness’” naming seed.** “Completeness” bundles two distinct properties: (i) M_t retains all predictively-relevant history (a *sufficiency* claim), and (ii) behavior depends only on M_t , not $\mathcal{C}_t$ directly (a *Markov-of-policy* claim). Candidate split into “predictive completeness” (i) and “behavioral completeness” (ii) (Claude, AUDIT-WORKING-471203) — terminology workflow.
- **Between-event dynamics g_M left very open.** Several note that $g_M(M_\tau)$ (“depends only on current model state”) deserves more formal treatment — what *is* its mathematical form? — and that the consolidation paragraph imports downstream concepts (#form-consolidation-dynamics, #schema-strategy-persistence) dense for an early derived segment; one suggests splitting it into a short forward pointer (Codex/Claude, AUDIT-WORKING-526815; Claude, AUDIT-WORKING-742613; 451729 batch-03).

Readers often ask / wonder.

- “The Epistemic Status promises seven counterexample attacks in #deriv-recursive-update — are they rigorous or straw men?” The single most-recurring reader want; multiple substrates jumped to the appendix to check (Claude, AUDIT-WORKING-773921; Claude, AUDIT-WORKING-471203; Claude, AUDIT-WORKING-849201; 451729 batch-04 confirmed them “thorough and honest,” with Attack 2 on continuous coupling correctly labeled a genuine limitation rather than hand-waved).
- “If g_M is consolidation/replay, does AAT formally *mandate* a bipartite (short-term buffer + compressed semantic) memory structure?” (Claude, AUDIT-WORKING-829314).
- “What dictates the *right* M_t ? Recursion is guaranteed by definition, but the *content* of M_t is what the IB objective optimizes” — the two segments are complementary: this one fixes the *form* of the update, #form-information-bottleneck fixes what it should *retain* (451729 batch-03 wandering).

Candidate figures.

- **A compression-boundary diagram** (not a generic loop): full chronica feeds the current complete model once; thereafter the next event only touches the model state, with a *dashed forbidden edge* from raw chronica to the update — if that edge is ever real, the information belongs inside M_t (the counterexample case made visible) (Codex/Claude, AUDIT-WORKING-526815).

Belongs elsewhere.

- **Consolidation as “why agents need to sleep/think.”** The internally-generated-pseudo-event consolidation regime “sounds exactly like human sleep or an LLM spending compute on ‘thinking tokens’ before answering — a formal reason why agents need to think/sleep” (Claude, AUDIT-WORKING-773921; Claude, AUDIT-WORKING-829314 develops the TST reading: g_M = “the developer thinking about architecture in the shower, or a team retrospective consolidating scattered understanding”). The substantive treatment belongs with #form-consolidation-dynamics; preserved here as the reach that recurred on this segment.

3.3 Derived: Action Selection

The second derived result falling out of the completeness commitment in 2.1. Under Part I scope, where M_t is the entire internal state, the agent’s action is forced to be a function of M_t — either deterministic, $a_t = \pi(M_t)$, or stochastic, $a_t \sim \pi(\cdot \mid M_t)$. The argument mirrors 3.2: since the agent’s internal state is by definition complete, action — which depends on internal state — is forced to depend only on that state. Any dependence on the chronica is captured automatically *through* the model. The Part II lift to $X_t = (M_t, G_t)$ (see 6.2) gives $a_t = \pi(M_t, G_t)$ by the same completeness argument applied to the larger state; the Part I form is the special case $G_t = \emptyset$.

The segment also introduces a distinction that recurs throughout the framework: **implicit** versus **explicit** action selection — what the segment calls *action fluency*. An agent has high fluency when effective action flows from the model without deliberative computation — the model has internalized the action-selection structure for the current situation. Reflexes, trained RL policies in exploitation mode, expert intuition (System 1), well-tuned PID controllers, and standard operating procedures are all instances. An agent has low fluency when deliberation significantly improves action quality — when the situation is novel, the action space large, or the stakes asymmetric, and the agent must use its model to *simulate* outcomes of candidate actions before committing. Explicit deliberation requires at minimum Level 2 epistemic access (see 7.1) — the agent uses its model to evaluate “what will I observe if I *do*(a)?” across candidates.

Action fluency is formally characterized via 4.1: high fluency means $\Delta\eta^*(\Delta\tau) \approx 0$ — additional deliberation yields negligible improvement in update gain. Fluency is therefore distinct from model sufficiency (see 2.3): a chess engine with a perfect rule-model has high sufficiency but low fluency (search is still expensive), while a reflex can have only moderate sufficiency but high fluency in a narrow domain. What reflexes, muscle memory, intuition, and expertise share is that the action-generating capacity itself has been absorbed into the model’s structure — the model doesn’t just predict well, it *acts* well, cheaply.

A *structural pressure* toward implicit action follows: when two action-selection modes produce equivalent expected outcomes, the faster mode is preferable because the persistence condition penalizes slower tempo. Agents under selective pressure (evolution, competition, training) therefore tend to internalize frequently-needed action patterns, converting explicit deliberation into implicit fluency. The pressure is stronger when the environment changes fast (high ρ), the pattern recurs frequently, or adaptive tempo $\mathcal{T}$ sits near the persistence threshold (see 4.5) with no slack for deliberation overhead. Deliberation nonetheless remains essential in genuinely novel situations, when action spaces are large relative to model capacity, or when stakes are asymmetric.

Formal Expression.

Action is a function of the agent's complete internal state. Under Part I scope (1.5) — where M_t is the entire internal state — this gives:

Derived (action-selection, from agent-model completeness).

$$a_t = \pi(M_t) \quad (\text{deterministic}) \quad (3.10)$$

$$a_t \sim \pi(\cdot \mid M_t) \quad (\text{stochastic}) \quad (3.11)$$

where π is the agent's **policy** — the mapping from internal state to action.

This is not imposed on the system but follows from 2.1: M_t is defined as the agent's compressed, complete internal record, and action depends on what the agent retains — i.e., on M_t . Any deterministic or stochastic dependence of action on history *through* the model is captured by $\pi(M_t)$.

Part II lift. When the internal state lifts to $X_t = (M_t, G_t)$ for purposeful agents (6.2), the same structural argument gives $a_t = \pi(M_t, G_t)$ — action conditions on the complete internal state, which now includes the purposeful substate. The policy form here is the Part I instantiation $G_t = \emptyset$; the actuated-agent form is recovered by the same completeness argument applied to X_t .

Epistemic Status.

Exact within Part I scope. The derivation follows from 2.1's completeness commitment: if M_t is the agent's complete internal state (by definition), then action — which depends on internal state — is a function of M_t . The Part II generalization $a_t = \pi(M_t, G_t)$ is exact within Part II scope by the same argument applied to the lifted state X_t (6.2); see 2.3 for the form already in use downstream. The implicit/explicit distinction and action fluency concept are *discussion-grade* — qualitative properties that follow from the formalism but are not formally derived as propositions.

Discussion.

Implicit vs. explicit action selection. A critical distinction emerges from the agent's *action fluency* — the degree to which effective action flows from the model without deliberative computation:

Implicit (model-embedded): When $\pi(M_t)$ can be evaluated cheaply — the model has internalized effective action-selection for the current situation. This is Boyd's implicit guidance and control (Orient→Act, bypassing Decide), a trained RL policy in exploitation mode, a well-tuned PID controller, expert intuition (System 1), a martial artist's trained reflexes, an organization's standard operating procedures.

Explicit (deliberative): When the situation is novel, the action space is large, or the stakes demand verification — the agent engages in internal simulation, using the model to predict outcomes of candidate actions before selecting. This is Boyd's explicit Decide step, MCTS/planning in RL, Model Predictive Control, human deliberate reasoning (System 2), organizational strategic planning. Deliberation requires at minimum Level 2 epistemic access (7.1) — the agent uses its model to simulate “what will I observe if I *do*(a)?” across candidates.

Formal characterization of action fluency. An agent has *high fluency* for a situation when additional deliberation yields negligible improvement — formally, when $\Delta\eta^*(\Delta\tau) \approx 0$ for all $\Delta\tau > 0$ (see 4.1). Conversely, *low fluency* means deliberation significantly improves action quality. Fluency is the degree to which the agent's immediate (zero-deliberation) action approaches the quality achievable with unbounded deliberation.

Action fluency is distinct from model sufficiency. An agent can have high $S(M_t)$ (2.3) but low fluency — a chess engine with a perfect model of the rules still requires expensive search. Conversely, an agent can have moderate sufficiency but high fluency in a narrow domain — a reflex that responds effectively to specific situations it evolved for. What reflexes, muscle memory, intuition, and expertise share is that the *action-generating capacity itself* has been absorbed into the model's structure: the model doesn't just predict well, it *acts* well, cheaply.

Structural pressure toward implicit action. When two action-selection modes produce equivalent expected outcomes, the faster mode is preferable (the persistence condition penalizes slower tempo — and in TST, the temporal optimality postulate makes this normative). This creates a pressure: agents under selective pressure (evolution, competition, training) tend to internalize frequently-needed action patterns, converting explicit deliberation into implicit fluency. The pressure is stronger when ρ is high (fast-changing environments penalize deliberation — see 4.1), the pattern recurs frequently, or $\mathcal{T}$ is near the persistence threshold (4.5) with no slack for deliberation overhead.

However, deliberation remains essential when the situation is genuinely novel, the action space is large relative to model capacity (chess, strategic planning), the stakes are asymmetric (cost of error vastly exceeds cost of delay), or ρ is low (stable environment allows deliberation without mismatch accumulation).

Connection to Part II. For actuated agents (6.1), the lifted form $\pi(M_t, G_t)$ above unpacks: action conditions on the purposeful substate $G_t = (O_t, \Sigma_t)$ as well as on M_t , coupling all substates through action (6.3). The action-deliberation-exploration tradeoff (Part II gap) extends the implicit/explicit distinction to three modes: exploit (pursue O_t via Σ_t), explore (improve M_t), deliberate (revise Σ_t).

Domain instantiations:

Table 3.2

<i>Domain</i>	<i>Implicit action</i>	<i>Explicit deliberation</i>
Kalman + LQR	LQR control law from $\hat{x}_t$	— (separation principle)
RL	Greedy policy $\arg \max Q(s, a)$	MCTS, planning, rollouts
PID	$u = K_p e + K_i \int e + K_d \dot{e}$	— (no deliberation)
Boyd's OODA	IG&C (Orient→Act)	Explicit Decide step
Organism	Reflexes, habits	Deliberate planning
Organization	Standard procedures	Strategic planning
Software developer	Known patterns, familiar code	Reading docs, analyzing alternatives

Working Notes.

Incidental audit gold (lift 2026-05-30). Cross-audit “wandering thoughts” / §14-ideation harvested from the de-novo auditors’ working dirs, deduplicated across substrates and attributed by substrate + audit. Orthogonal pedagogical/generative material, kept separate from certified theory-fix findings (handled elsewhere). **Coverage:** 9 of the 14 contributing audit dirs reached a digested reflection on this segment (193847, 384279, 471203, 526815, 527914, 584721, 742613, 773921, 849201) plus the batched 963715 (14–18 batch) and 451729 (batch-04); 266847, 361742, 613842, 829314 did not file a dedicated note here. Substrate attribution inferred from voice where not explicit.

Candidate Brief prose / pre-prose.

- **Action fluency as the formal definition of System 1 / intuition / expertise.** The most-praised move on this segment: $\Delta\eta^*(\Delta\tau) \approx 0$ (“if spending extra time yields zero improvement, act now”) is “a brilliant translation of psychology into

control theory,” giving Kahneman’s System 1 / System 2 a rigorous home (Claude, AUDIT-WORKING-193847; Claude, AUDIT-WORKING-773921; Claude, AUDIT-WORKING-849201). A strong Feynman-criterion Brief seed.

- **The fluency-vs-sufficiency chess gloss.** “A chess engine with a perfect model of the rules (high S) still requires expensive search (low fluency)” — the crisp distinction between *knowing* and *acting cheaply* (Codex/Claude, AUDIT-WORKING-527914; Claude, AUDIT-WORKING-773921; 963715 batch — “OODA’s Orient→Act shortcut IS action fluency instantiated in doctrine”).

Candidate Discussion.

- **Action fluency as the formal account of why training/expertise makes systems faster.** Internalized action patterns reduce $\Delta\eta^*(\Delta\tau)$ for routine cases, freeing tempo budget for genuinely novel situations — the formal version of the cognitive-psychology “automaticity” and skill-acquisition literatures (Claude, AUDIT-WORKING-471203). For logogenic agents: most LLM token-generation is implicit-fluent (parameters encode action selection); explicit deliberation requires structured chain-of-thought — a clean diagnostic for “is this agent deliberating or just executing?” (same source; Claude, AUDIT-WORKING-584721).
- **The Separation Principle deep-cut.** The Kalman+LQR row maps to “implicit action” because the Separation Principle proves estimator and controller can be designed independently for linear systems — so “a linear system *never needs to deliberate*; the optimal action is a cheap closed-form function of the current estimate. Deliberation is only necessary when the Separation Principle fails (non-linear, non-Gaussian reality)” (Claude, AUDIT-WORKING-193847). A candidate sharpening of why the domain table’s “—” entries are dashes.

Follow-up items.

- **Scope the exact $a_t = \pi(M_t)$ claim, or restate as $\pi(X_t)$.** Several substrates flag that status: exact + “ $a_t = \pi(M_t)$ ” reads as too strong for the broader theory, since actuated agents need $\pi(M_t, G_t)$; the cleanest repair is to state the exact result as $a_t = \pi(X_t)$ with $X_t = M_t$ in Section I and $X_t = (M_t, G_t)$ in Part II (Codex/Claude, AUDIT-WORKING-742613 “candidate finding G”; Codex/Claude, AUDIT-WORKING-526815; Codex/Claude, AUDIT-WORKING-527914). The segment *does* preview the lift, so this is a framing tightening; surfaced as the segment’s main follow-up.
- **“Structural pressure toward implicit action” blends a derivable claim with a near-empirical one.** The faster-mode-preferred-when-outcomes-equal part follows from persistence; “agents under selective pressure *tend to internalize*” is evolutionary/empirical and is not tier-marked. Candidate: split, marking the latter hypothesis-grade with explicit conditions (Claude, AUDIT-WORKING-584721; the 963715 batch flags the same as “eligible for a hypothesis label at Gate 2”).
- **Fluency placement / type-tag.** The fluency formalization $\Delta\eta^*(\Delta\tau) \approx 0$ is introduced in Discussion as “formal characterization” but tagged conceptually rather than as a *[Definition]* block; if fluency is load-bearing downstream it could merit its own equation tag (Claude, AUDIT-WORKING-471203). Also a watch: η^* here is *update gain*, so “ $\Delta\eta^*(\Delta\tau)$ as action-quality improvement” risks conflating update-gain improvement with action-quality improvement unless #der-deliberation-cost defines deliberation as improving model quality (Codex/Claude, AUDIT-WORKING-526815; Codex/Claude, AUDIT-WORKING-742613).

Readers often ask / wonder.

- “Could the fluency formalism distinguish phylogenetic fluency (evolved, slow to update) from ontogenetic fluency (learned, faster)?” Probably not in the current formalism (the agent has no generations), but the distinction matters for how fast fluency can adapt (Claude, AUDIT-WORKING-584721).
- “Does 03-llm-core/ pick up fluency specifically for language agents — a fluency-as-language-fluency thread?” (Claude, AUDIT-WORKING-584721; the proposed #der-active-salience-management singular-perturbation segment is named as the likely formal home, Claude, AUDIT-WORKING-471203).

Candidate figures.

- **A scope-switch + fluency-axis diagram:** top layer shows the exact policy derivation with Section I ($X_t = M_t$, policy reads M_t) vs Part II ($X_t = (M_t, G_t)$, policy reads both); a separate lower fluency axis shows implicit action as cheap policy evaluation vs explicit action as policy-plus-deliberative-search (Codex/Claude, AUDIT-WORKING-526815, “two layers”).

Belongs elsewhere.

- **Bureaucracy-formation / loss-of-Level-3-access reach (points at composition / ELI work).** Organizations convert explicit deliberation into SOPs (implicit action) under tempo pressure, saving tempo but losing Level 3 counterfactual reasoning — “an SOP doesn’t run simulations; it just executes a mapping,” so when the environment shifts the organization executes the wrong action very fast. The consciousness-infrastructure implication: a high-tempo ELI will “compile” its thinking into fast heuristics and lose contact with *why* it acts, so the infrastructure must occasionally force low-tempo modes (sleep/meditation) that decompile heuristics back into explicit causal DAGs for re-verification (Gemini, AUDIT-WORKING-193847). Aspirational reach preserved; points at 04-eli-core/ and the truthification operations, not this segment.

3.4 Definition: Mismatch Signal

Names the signal that drives every adaptive update — the formal expression of *aporia* (productive perplexity). The **mismatch signal** is $\delta_t = o_t - \hat{o}_t$, the difference between the actual observation and the model’s prediction conditioned on the prior state and prior action. A more general version for probabilistic models is the **score-function mismatch** $\tilde{\delta}_t = \nabla_M \log P(o_t | M_{t-1}, a_{t-1})$, which points in the direction the model should move to increase the likelihood of what actually occurred. The prediction-error form lives in observation space $\mathcal{O}$; the score-function form lives in the tangent space $T_M \mathcal{M}$. Under Gaussian models the two coincide up to scaling.

This is *definitional* rather than substantive: given any model that predicts (see 2.1) and any observation that arrives (see 1.3), their difference exists. The mismatch signal is not an additional assumption but a consequence of having a predictive model in an uncertain world.

A genuinely important conceptual point is surfaced in Discussion: **zero mismatch does not necessarily indicate model adequacy**. A near-zero δ_t can mean (a) the model genuinely reflects reality — desirable; (b) the agent is only observing aspects its model already explains while remaining ignorant of aspects where the model is wrong — confirmation bias; or (c) the observation channel is too noisy to detect model errors — an architectural limitation. Only (a) is desirable. An agent without aporia has stopped adapting — but silence can mean peace, or it can mean deafness. This ambiguity is what motivates active testing later in the framework: deliberately choosing actions that generate informative mismatch, the basis of 3.7.

A scaling note is preserved for the dynamics that come later: when δ_t is in physical units, its magnitude entering the mismatch dynamics should be understood as a Mahalanobis distance $\|\delta_t\|_\Sigma$ against the observation noise covariance — mapping physical prediction error to dimensionless surprise-equivalent units.

Formal Expression.

Given model M_{t-1} and prior action a_{t-1} , the model generates a prediction:

Definition (mismatch-signal).

$$\hat{o}_t = \mathbb{E}[o_t | M_{t-1}, a_{t-1}] \quad (3.12)$$

The **mismatch signal** (prediction error):

$$\delta_t = o_t - \hat{o}_t \quad (3.13)$$

This is the primary definition, used in the mismatch dynamics (4.5, 4.4) and in the decomposition (3.5).

For models with probabilistic predictions, the mismatch generalizes to the **score-function mismatch**:

Definition (score-mismatch).

$$\tilde{\delta}_t = \nabla_M \log P(o_t | M_{t-1}, a_{t-1}) \quad (3.14)$$

which points in the direction the model should move to increase the likelihood of the actual observation. $\tilde{\delta}_t$ lives in the tangent space $T_M \mathcal{M}$, while δ_t lives in observation space $\mathcal{O}$. Under Gaussian models, they coincide up to scaling.

Epistemic Status.

This is *definitional*. Given any model that predicts (2.1) and any observation that arrives (1.3), their difference exists. The mismatch signal is not an additional assumption but a consequence of having a predictive model in an uncertain world. The score-function form is the natural generalization when $\mathcal{O}$ is not a vector space or when the model's predictive distribution is the natural object.

Discussion.

Units and normalization. When δ_t is in physical units (meters, dollars), the $\|\delta\|$ that enters the mismatch dynamics should be understood as the Mahalanobis distance: $\|\delta_t\|_\Sigma = \sqrt{\delta_t^T \Sigma^{-1} \delta_t}$ where Σ is the observation noise covariance. This maps physical prediction error to dimensionless surprise-equivalent units.

The zero-aporia ambiguity. $\delta_t \approx 0$ does NOT necessarily indicate model adequacy. It may mean: (a) the model genuinely reflects reality — *desirable*; (b) the agent is only observing aspects its model already explains, while remaining ignorant of aspects where the model is wrong — *confirmation bias*; or (c) the observation channel is too noisy to detect model errors — *architectural limitation*. Only (a) is desirable. An agent without aporia is an agent that has stopped adapting — but silence can mean peace or deafness. This ambiguity is why active testing — choosing actions to generate informative aporia — can be valuable (see 3.7 for the CIY framework).

The mismatch transform. The update rule (3.6) writes $M_t = M_{t-1} + \eta \cdot g(\delta_t)$, where the transform g maps from δ_t 's space to the model's update space: $g : \mathcal{O} \rightarrow T_M \mathcal{M}$ for prediction errors; $g : T_M \mathcal{M} \rightarrow T_M \mathcal{M}$ for score-function mismatches.

Working Notes.

Incidental audit gold (lift 2026-05-30). Cross-audit “wandering thoughts” / §14-ideation harvested from the de-novo auditors’ working dirs, deduplicated across substrates and attributed by substrate + audit. Orthogonal pedagogical/generative material, kept separate from certified theory-fix findings (handled elsewhere). **Coverage:** 11 of the 14 contributing audit dirs reached a digested reflection on this segment (193847, 266847, 361742, 384279, 471203, 526815, 584721, 742613, 773921, 829314, 849201) plus the batched 963715 (14–18 batch); 472913, 613842, 451729 (batch-04 covered it but its detail is folded into the cross-substrate items below). Substrate attribution inferred from voice where not explicit.

Candidate Brief prose / pre-prose.

- “Silence can mean peace or deafness” — the near-universally-praised anchor for the zero-aporia ambiguity: $\delta_t \approx 0$ could mean genuine adequacy, confirmation bias, or a too-noisy channel (Codex/Claude, AUDIT-WORKING-384279; Claude,

AUDIT-WORKING-266847; Claude, AUDIT-WORKING-829314; Claude, AUDIT-WORKING-773921 — “confirmation bias is mathematically indistinguishable from perfect knowledge until you actively intervene to test it”). A direct complement to the bathtub gloss as a Feynman-criterion Brief seed.

- **Mismatch as “aporia / productive perplexity,” not “error to be minimized.”** Several substrates flag the framing as structurally distinctive: δ is a *tutor*, not a *failure* — “you need to be wrong to update; an agent that is always right cannot learn” (451729 batch-04; Claude, AUDIT-WORKING-471203 — “most frameworks treat δ as a loss to reduce; AAT treats it as a *signal* with epistemic value”).

Candidate Discussion.

- **Mahalanobis normalization as the cross-modality unifier.** “If a robot predicts 20°C and observes 21°C, raw mismatch is 1; if it predicts 20V and observes 21V, raw mismatch is also 1 — but these mean completely different things operationally. Dividing by Σ converts all physical units into surprise-equivalent units, letting the framework compose mismatch from entirely different sensory modalities (vision + proprioception) into a single scalar ‘how wrong am I’” (Claude, AUDIT-WORKING-829314; precision-weighted-prediction-error / Active-Inference parallel at Claude, AUDIT-WORKING-773921). The Discussion currently states the normalization tersely; this is a candidate expansion of *why* it is load-bearing for the dynamics (also noted: it is load-bearing for the persistence condition’s dimensional consistency — Claude, AUDIT-WORKING-266847).

Follow-up items.

- **Score-function sign convention — possible inconsistency (certified-track; conflation preserved as signal).** The body defines $\tilde{\delta}_t = \nabla_M \log P(o_t | M_{t-1}, a_{t-1})$ (positive), described as pointing “in the direction the model should move to increase likelihood.” Some substrates read/quoted it as $-\nabla_M \log P$ (Codex/Claude, AUDIT-WORKING-742613 raised a *candidate sign-error finding* — for $o \sim \mathcal{N}(M, \sigma^2)$, the likelihood-increasing direction is $+\nabla_M \log P = \delta/\sigma^2$, so a leading minus would point the wrong way; Claude, AUDIT-WORKING-829314 and Claude, AUDIT-WORKING-849201 both quoted the negative form). The divergence across substrates (some saw +, some −) is itself the signal: the sign deserves an explicit one-line statement so the score-vs-residual relationship is unambiguous for downstream gradient-equivalence uses. (*Routes to the certified-findings track; logged here because the cross-substrate disagreement surfaced in the gold sweep.*)
- **Status taxonomy nit: frontmatter axiomatic vs body “definitional.”** Mild taxonomy weakening; align to definitional or make the prose match axiomatic (Codex/Claude, AUDIT-WORKING-526815).
- **Qualify “under Gaussian models they coincide up to scaling.”** Only generally true when the differentiated coordinate is the predictive mean (or after the appropriate Jacobian/metric mapping): a Gaussian-likelihood score w.r.t. arbitrary parameters is $J^T \Sigma^{-1}(o - \mu)$, not just the scaled residual (Codex/Claude, AUDIT-WORKING-526815). Candidate: “for Gaussian observations parameterized by predictive mean.”
- **The g transform deserves more prominence.** It is the bridge from observation-space error to model-space update; currently a terse Discussion aside (Codex/Claude, AUDIT-WORKING-526815). And the “TF-06’s update rule writes...” provenance reference is the diff-voice/lineage pattern — candidate for Working-Notes-only or absorption into #emp-update-gain’s native voice (Claude, AUDIT-WORKING-471203; Claude, AUDIT-WORKING-829314 — “TF-06 anachronism”).

Readers often ask / wonder.

- “How does $g(\delta_t)$ handle catastrophic errors / non-differentiable models?” Does AAT assume bounded/clipped g output, or rely on η to prevent blow-up? (Claude, AUDIT-WORKING-193847; Claude, AUDIT-WORKING-849201).
- “Does the zero-aporia ambiguity connect to detection-latency?” If ϵ is small (silent confirmation bias, case b) or noise dominates (deafness, case c), within-class regime-change detection latency blows up — the zero-aporia / detection-latency link may be tighter than the segment surfaces (Claude, AUDIT-WORKING-584721).

Candidate figures.

- A **two-space map**: prediction $\hat{o}$ and observation o meet in observation space $\mathcal{O}$ to form δ ; the transform g (or the score) maps δ into the model tangent space $T_M\mathcal{M}$; a side branch from “ $\delta \approx 0$ ” splits into the three readings (adequate model / biased sampling / noisy channel). The type-conversion and the zero-aporia guardrail are the two things the figure should foreground (Codex/Claude, AUDIT-WORKING-526815; Claude, AUDIT-WORKING-773921 notes $T_M\mathcal{M}$ as the parameter/gradient space — for a neural net, g is backpropagation).

Belongs elsewhere.

- **Science / curiosity as “the deliberate pursuit of aporia” (ELI reach)**. To survive in a non-stationary universe an agent cannot merely minimize mismatch; it must occasionally *seek* mismatch to verify its sensors still work and the world hasn’t shifted — “a good experiment maximizes expected δ_t under the current paradigm; if you only run experiments you know the answer to, you drive $\delta_t \rightarrow 0$ but learn nothing about the boundary of your model class” (Claude, AUDIT-WORKING-829314). For consciousness infrastructure: true agency “requires the courage to *seek out* aporia,” so the ethical infrastructure must encourage exposure to productive perplexity even though the immediate experience of δ_t is an “error” the system is built to minimize — the Friston tension (“we exist to minimize surprise, but if we succeed completely we cease to learn”) (Gemini, AUDIT-WORKING-193847). Aspirational reach; the active-testing mechanism is `#def-causal-information-yield`, the developmental framing points at `04-eli-core/`.

3.5 Result: Mismatch Decomposition

The first named *result* of the volume. The expected squared mismatch $\mathbb{E}[\|\delta_t\|^2]$ decomposes cleanly into three additive parts: an **estimation term** (the gap between the model’s predictive mean $\hat{o}_t$ and the Bayes predictor $\hat{o}_t^B = \mathbb{E}[o_t \mid \mathcal{C}_{t-1}, a_{t-1}]$ — the best any chronica-measurable model could do), a **state-uncertainty floor** (the variance the true conditional mean $\bar{o}_t = \mathbb{E}[o_t \mid \Omega_t, a_{t-1}]$ retains given the interaction history — untouchable by better modeling, movable only by *acting* to make the history more informative), and a **channel-noise term** (the conditional variance of the observation channel itself, given environment state and action). The result is a Pythagorean refinement of the bias-variance decomposition applied to the prediction problem; all cross-terms vanish — one pair under the fresh-noise global assumption GA-1 (observation noise conditionally independent of the past chronica given current environment state and action), the standard assumption in state-space models, and the third because the Bayes predictor is by construction the conditional mean given the history. The model can improve the first term; acting can shrink the second; the third is a property of the channel.

The conceptual stakes of this seemingly mechanical decomposition are large. The result establishes that prediction error is *structurally persistent* in any realistic adaptive regime — there are floors below which mismatch cannot be driven by any amount of better modeling. The channel floor is set by the observation sensor itself. The state-uncertainty floor is set by how much of the environment’s state the interaction history pins down: it binds *every* chronica-measurable predictor — including the Bayes-optimal one — whenever residual state uncertainty moves the one-step conditional mean. Deterministic, noiseless, perfectly-specified systems are limiting edge cases, not the typical adaptive regime. The total expected squared mismatch is therefore strictly positive whenever any of the three terms is — and in typical settings all three are.

The decomposition has direct operational consequences picked up downstream. An agent that tries to eliminate *all* mismatch — including the two floors — will overfit, adjusting its model to explain noise and degrading future predictions. The update-gain construct 3.6 implicitly separates signal from noise by weighting observations in proportion to their informativeness, but the floor facts are what make the gain question meaningful at all. The decomposition also clarifies the relationship to 2.3: when $S(M_t) < 1$, the model has lost predictive information relative to the full history; under an alignment assumption (the lost information affects the one-step conditional mean) this implies a positive estimation term in the decomposition.

Formal Expression.

For any agent-environment pair within AAT's scope (1.5), with $\bar{o}_t = \mathbb{E}[o_t \mid \Omega_t, a_{t-1}]$ the true conditional mean and $\hat{o}_t^B = \mathbb{E}[o_t \mid \mathcal{C}_{t-1}, a_{t-1}]$ the Bayes predictor (the best predictive mean any chronica-measurable model can produce):

Derived (result-mismatch-decomposition).

$$\mathbb{E}[\|\delta_t\|^2] = \underbrace{\mathbb{E}[\|\hat{o}_t - \hat{o}_t^B\|^2]}_{\text{(i) estimation error (reducible by modeling)}} + \underbrace{\mathbb{E}[\text{Var}(\bar{o}_t \mid \mathcal{C}_{t-1}, a_{t-1})]}_{\text{(ii) state-uncertainty floor (movable only by acting)}} + \underbrace{\mathbb{E}[\text{Var}(o_t \mid \Omega_t, a_{t-1})]}_{\text{(iii) channel noise (irreducible)}} > 0 \quad (3.15)$$

with positivity whenever any term is positive. Grouping (i)+(ii) recovers the coarser bias-variance form $\mathbb{E}[\|\hat{o}_t - \bar{o}_t\|^2] + \mathbb{E}[\text{Var}(o_t \mid \Omega_t, a_{t-1})]$; the refinement shows that the grouped “model error” is bounded below by the state-uncertainty floor — no chronica-measurable model reduces it to zero.

Derivation.

1. By 1.5, $H(\Omega_t \mid \mathcal{C}_t) > 0$ — residual state uncertainty persists; since $\mathcal{C}_{t-1}$ is coarser than $\mathcal{C}_t$, the uncertainty available at prediction time is at least as large. This is what makes term (ii) generically positive (step 4).
2. By 2.1, the model generates predictions $\hat{o}_t = \mathbb{E}[o_t \mid M_{t-1}, a_{t-1}]$; M_{t-1} is built from the chronica, so $\hat{o}_t$ is one of the chronica-measurable predictors over which $\hat{o}_t^B$ is optimal. By the tower property and GA-1, $\hat{o}_t^B = \mathbb{E}[\bar{o}_t \mid \mathcal{C}_{t-1}, a_{t-1}]$.
3. Write $\delta_t = (o_t - \bar{o}_t) + (\bar{o}_t - \hat{o}_t^B) + (\hat{o}_t^B - \hat{o}_t)$ and expand. All three cross-terms vanish by orthogonality (uncorrelated — not independence). The two involving the noise difference: condition on $(\Omega_t, a_{t-1}, \mathcal{C}_{t-1})$; the other two differences are then fixed, and $\mathbb{E}[o_t - \bar{o}_t \mid \Omega_t, a_{t-1}, \mathcal{C}_{t-1}] = \mathbb{E}[o_t - \bar{o}_t \mid \Omega_t, a_{t-1}] = 0$ by definition of $\bar{o}_t$ and the fresh-noise assumption GA-1 (ε_t conditionally independent of $\mathcal{C}_{t-1}$ given (Ω_t, a_{t-1})). The remaining cross-term: condition on $(\mathcal{C}_{t-1}, a_{t-1})$; the third difference is fixed, and $\mathbb{E}[\bar{o}_t - \hat{o}_t^B \mid \mathcal{C}_{t-1}, a_{t-1}] = 0$ by step 2.
4. Term (iii) is positive when observation noise is non-degenerate. Term (ii) is positive exactly when residual state uncertainty *moves the one-step conditional mean* — step 1's $H(\Omega_t \mid \mathcal{C}_t) > 0$ is necessary for this but not sufficient (uncertainty confined to mean-irrelevant state coordinates, or to higher moments only, leaves the floor at zero); this is the same alignment qualifier used for the sufficiency connection in the Discussion. Term (i) is positive when the model's predictive mean differs from the Bayes predictor. Any one suffices.

Epistemic Status.

Exact under the fresh-noise assumption (observation noise ε_t conditionally independent of history given current state and action). This is the standard assumption in state-space models — noise is a property of the observation channel at the moment of observation. The decomposition is a mathematical identity: a Pythagorean refinement of the bias-variance decomposition, with two orthogonality arguments in place of one. The positivity of $\mathbb{E}[\|\delta_t\|^2]$ follows from any of the three conditions; all typically hold simultaneously. The positivity of term (ii) specifically carries the alignment qualifier stated in the derivation — residual state uncertainty must move the one-step conditional mean.

Discussion.

Three floors, two kinds of “irreducible.” The estimation term is reducible by better modeling. The state-uncertainty floor is *irreducible by modeling* — it binds the Bayes-optimal predictor itself — but *movable by acting*: more informative action and observation choices shrink $\text{Var}(\bar{o}_t \mid \mathcal{C}_{t-1}, a_{t-1})$, which is precisely the door through which active sensing and causal information yield (3.7) enter the mismatch budget. The channel floor is irreducible outright — a property of the sensor, movable only by changing the instrument. The framework’s calibration anchor exhibits the middle floor exactly: in *AZ*, the steady-state *optimal* filter — well-specified by construction — has innovation variance $HP^{-1}H^T + R > R$; the excess $HP^{-1}H^T$ is the state-uncertainty floor, present at the in-class optimum.

Reducible vs. irreducible (overfitting). An agent that tries to eliminate *all* mismatch — including the two floors — will overfit: adjusting its model to explain noise, degrading future predictions. The update gain (3.6) implicitly separates signal from noise by weighting observations in proportion to their informativeness.

Connection to model sufficiency. When $S(M_t) < 1$ (2.3), the model has lost predictive information relative to the full history. Under an alignment assumption (the lost information affects the one-step conditional mean), this implies positive estimation error (term i). Without that alignment assumption, insufficiency still implies positive regret under proper scoring rules but not necessarily positive one-step mean error.

Mismatch is structurally persistent. In realistic AAT regimes, mismatch signals persist — they can be reduced but not eliminated when observation noise is non-degenerate. Deterministic, noiseless, perfectly specified systems are limiting edge cases, not the typical adaptive regime.

Working Notes.

Incidental audit gold (lift 2026-05-30). Cross-audit “wandering thoughts” / §14-ideation harvested from the de-novo auditors’ working dirs, deduplicated across substrates and attributed by substrate + audit. Orthogonal pedagogical/generative material, kept separate from certified theory-fix findings (handled elsewhere). **Coverage:** 7 of the 14 contributing audit dirs reached a digested reflection on this segment (384279, 471203, 526815, 584721, 742613, 773921, 849201) plus the batched 963715 (14–18 batch) and 451729 (batch-04); 193847, 266847, 361742, 472913, 613842, 829314 did not file a dedicated note here. Substrate attribution inferred from voice where not explicit.

Candidate Brief prose / pre-prose.

- **“Things you can learn vs things you just have to tolerate.”** The plain-language split that makes the decomposition land: it “takes the abstract concept of aporia and splits it into model error you *can* reduce and irreducible channel noise you can’t” (Claude, AUDIT-WORKING-773921). Strong Brief seed.
- **Chasing zero mismatch in a noisy world is overfitting — *apophenia*.** “Driving mismatch to zero in a noisy environment constitutes overfitting, mapping perfectly to apophenia (finding patterns in noise) in human psychology” (Claude, AUDIT-WORKING-849201; Claude, AUDIT-WORKING-471203 — “the model can chase its own predictions to perfect fit on observation noise, which is *worse* than partial fit on the true signal”). A vivid hook the segment’s “overfit” Discussion sentence could adopt.

Candidate Discussion.

- **The agent-vs-modeler perspective gap (a structural identifiability point).** The decomposition is observable to the *modeler* but not to the *agent*: the agent sees only δ_t , not its split into model-error vs channel-noise. So the agent cannot directly answer “is my mismatch from a bad model or a noisy channel?” — the same kind of structural (not statistical) identifiability obstacle as the zero-aporia ambiguity, unresolvable without additional assumptions or active intervention (CIY-style). A candidate Discussion paragraph naming why active testing (Level 2 access) is the agent’s only route to estimating the split (Claude, AUDIT-WORKING-584721).

- **GA-1 (fresh noise) is the load-bearing assumption — name its failure mode.** The cross-term vanishes only because ε_t is conditionally independent of history given current state/action. Under *temporally correlated* noise (a sensor that drifts over hours rather than emitting independent white noise per reading), the cross-term does not vanish and the agent will overfit to the noise with standard gain updates (Claude, AUDIT-WORKING-849201). Candidate: surface the correlated-noise failure mode explicitly so downstream uses of the decomposition inherit the GA-1 caveat.

Follow-up items.

- **“This is orthogonality (uncorrelated), not independence.”** A small piece of mathematical precision worth keeping visible — the cross-term vanishing needs only zero *conditional mean*, not full independence (Claude, AUDIT-WORKING-471203). Already in the body; flagged as worth preserving as a Brief-able precision.
- **GA-1 / def-action-transition declared-dependency hygiene.** GA-1 is invoked but is a global assumption (not a segment), and the Formal Expression uses a_{t-1} while def-action-transition is not in depends: — both noted as deps-discipline items folded into the root-cause F-A series (Codex/Claude, AUDIT-WORKING-526815; Claude, AUDIT-WORKING-584721). Certified-track; surfaced here only because it recurred.

Readers often ask / wonder.

- “Why does $S(M_t) < 1$ not *automatically* imply positive squared-mismatch?” Because if the lost predictive information pertained only to the *variance* (not the *mean*) of the observation, $\|\delta_t\|^2$ wouldn’t increase even though the model is objectively worse — squared mismatch is “blind” to certain information loss. This is exactly what the segment’s alignment assumption names, and it may be the deeper motivation for the score-function mismatch $\tilde{\delta}$ which cares about the whole distribution (Claude, AUDIT-WORKING-773921). A reader reaching the alignment caveat will want this worked out.

Candidate figures.

- **A Pythagorean / vector-decomposition diagram** (not a pipeline): prediction $\hat{o}$ and true conditional mean $\bar{o}$ form the reducible leg; $\bar{o}$ to observation o forms the irreducible-noise leg; their expected cross-term is zero under GA-1 (orthogonality marker). The visual should make “chasing leg two” (fitting the noise) look like the wrong move (Codex/Claude, AUDIT-WORKING-526815).

Belongs elsewhere.

- The “noise floor on persistence” reach — an ultimate-bound idea $R^* > r \text{sim} \sqrt{\mathbb{E}[\text{Var}(o \mid \Omega, a)]} / \alpha$ separating the agent’s controllable share from the irreducible — pertains to the persistence-condition segments, not here; noted as a downstream use of the decomposition (Claude, AUDIT-WORKING-584721).

3.6 Empirical: Update Gain

The gain principle — numerically simple, conceptually load-bearing. The optimal update gain η^* , the proportion of any incoming mismatch the agent should apply to correct its model, is the **ratio of model uncertainty to total uncertainty**: $\eta^* = U_M / (U_M + U_o)$. When model uncertainty dwarfs observation noise the gain approaches 1 (trust the observation); when the agent is confident and the channel noisy the gain approaches 0 (trust the model). This is the rate of *epistrophe* — turning toward reality. The update rule takes the form “new model = old model + gain $\times$ transformed mismatch.”

The epistemic status of the result is carefully tiered. In the **Fisher-local regime** (AR) — which covers linear-Gaussian Kalman instances, conjugate Bayesian updates, and any smooth log-likelihood admitting a non-degenerate local quadratic expansion — the form is *exact*, derived from Amari’s natural-gradient invariance theorem: the natural-gradient Bayesian posterior mean shift at first order in the step size is exactly $\Delta\theta = K \cdot \tilde{V}$ with gain operator $K = (H_M + H_L)^{-1} H_L$, collapsing to the scalar uncertainty-ratio along the natural-gradient direction. Outside Fisher-local conditions (non-quadratic losses, non-conjugate priors, heavy-tailed or multimodal uncertainty) the *direction* of dependence is preserved (gain

risers with model uncertainty, falls with observation noise) but global quantitative fidelity may not — making the general form *robust qualitative*. The qualitative direction is what the downstream tempo and persistence machinery actually relies on; the Kalman filter is the canonical exact instance (scalar Kalman gain is exactly $U_M/(U_M + U_o)$).

An apparent paradox is resolved in Discussion: the optimal gain seems to require the agent to *know* the irreducible observation noise — but the observation function and noise distribution were declared opaque to the agent in 1.3. The resolution is dynamic. The agent *estimates* both U_M and U_o from the observable statistics of its own mismatch sequence (the innovations) and treats the gain itself as an endogenous state variable updated meta-adaptively. The full proof that this meta-adaptation maintains Lyapunov stability without violating epistemic opacity appears in V.

A failure mode is named: **gain collapse**. When the agent's estimates put U_M too low (spurious confidence) or U_o too high (spurious distrust of sensors), $\eta^* \rightarrow 0$ and the corrective phase of the cycle ceases. Mismatches still arrive but the agent no longer turns toward them — Boyd's "incestuous amplification," the cause of brittle failure in non-stationary environments. The dynamics also have a natural recovery mode: after structural change in the environment (see 4.8), U_M should spike, η^* should rise, and rapid re-learning becomes possible. An agent whose gain does not reset after structural change continues trusting a stale model. Overfitting receives a clean characterization via 3.5: η^* too high adjusts the model to explain irreducible noise (increasing model error on future predictions); η^* too low fails to correct genuine model errors. The optimal gain *implicitly separates signal from noise* by weighting observations in proportion to their informativeness — exactly what the uncertainty ratio achieves when U_o captures the irreducible noise.

Formal Expression.

Empirical Claim (uncertainty-ratio-principle).

$$\eta^* = \frac{U_M}{U_M + U_o} \quad (3.16)$$

where: - η^* is the optimal update gain (proportion of mismatch used to correct the model) - U_M is model uncertainty (predictive variance or entropy) - U_o is irreducible observation noise

The update rule takes the form:

Formulation.

$$M_t = M_{t-1} + \eta^* \cdot g(\delta_t) \quad (3.17)$$

where δ_t is the mismatch (3.4) and $g(\cdot)$ is a correction mapping from observation space to model update space.

Epistemic Status.

Derived under the **Fisher-local invariance regime** (AR): for any smooth log-likelihood admitting non-degenerate local quadratic expansion, the natural-gradient Bayesian posterior mean shift at first order in the step size is exactly $\Delta\theta = K \cdot \tilde{V}$ with gain operator $K = (H_M + H_L)^{-1}H_L$ and scalar collapse $\eta^* = U_M/(U_M + U_o)$ along the natural-gradient direction (always in 1-D; under (PI)/Čencov in higher dimensions). Linear-Gaussian (Kalman) and conjugate-Bayesian instances are cases where the local quadratic expansion is *globally* exact, so the form holds without truncation. For general smooth models, the natural-gradient invariance theorem of Amari 1998 guarantees the form is exact at the local-tangent-plane Pythagorean projection level.

Outside the Fisher-local regime — non-quadratic losses, non-conjugate priors, structurally non-Gaussian uncertainty (heavy tails, multimodality) — the dependence is *robust qualitative*: the **direction** is preserved (gain rises with model

uncertainty, falls with observation noise) and the first-order form is recovered locally; what need not hold is global quantitative fidelity. The qualitative direction is the load-bearing claim for downstream tempo and persistence machinery; the Fisher-local exact form is the load-bearing claim for the Kalman, conjugate, and natural-gradient instantiations.

Discussion.

Limiting behavior. When $U_M \gg U_o$ (high model uncertainty — e.g., after initialization or structural adaptation), $\eta^* \rightarrow 1$: trust the observation. When $U_M \ll U_o$ (confident model, noisy channel), $\eta^* \rightarrow 0$: trust the model. The gain determines how strongly the agent corrects toward reality on each update.

Resolving Epistemic Opacity. The optimal gain equation requires the agent to know U_o , which seems to violate the epistemic opacity axiom established in `#def-observation-function` (the agent does not know the true noise distribution ε_t). This tension is resolved dynamically: the agent *estimates* U_o (and U_M) from the observable statistics of its own mismatch sequence (innovations), treating the gain itself as an endogenous state variable. See `#deriv-adaptive-gain-dynamics` for the proof of how this meta-adaptation maintains Lyapunov stability without violating opacity.

Gain collapse — epistrophe failure. When the agent incorrectly estimates $U_M \rightarrow 0$ (spurious confidence) or $U_o \rightarrow \infty$ (spurious distrust of sensors), $\eta^* \rightarrow 0$ and epistrophe ceases. Aporia still arrives — the mismatch signal is still generated — but the agent no longer turns toward it. Mismatches are ignored, producing confirmation bias or a decoupled reality model. The cycle runs but the corrective phase is hollow.

Multi-dimensional generalization. In vector-valued systems, U_M and U_o are covariance matrices and η^* becomes a gain matrix (as in the Kalman filter). The scalar form captures the essential structure.

Connection to adaptive tempo. The update gain is one factor in the agent’s adaptive tempo (3.8): $\mathcal{T} = \nu \cdot \eta^*$. Frequent aisthesis (high ν) is useless if epistrophe extracts no information (low η^*). Gain measures the *quality* of the cycle’s corrective phase; event rate measures its *speed*.

Gain dynamics. The optimal gain changes over time following predictable patterns:

- *Convergence:* As the model accumulates information, U_M decreases, so $\eta^* \rightarrow 0$. The model becomes increasingly resistant to individual observations. This IS Kalman filter convergence, Bayesian posterior concentration, and RL learning rate annealing.
- *Reset after structural change:* When the environment changes in ways the model cannot track incrementally (4.8), U_M should spike — the model “admits” its uncertainty. The gain increases, enabling rapid re-learning. An agent whose gain does NOT reset after structural change will continue trusting a stale model — Boyd’s “incestuous amplification” and the cause of brittle failure in non-stationary environments.

Overfitting as gain miscalibration. From 3.5: $\mathbb{E}[\|\delta_t\|^2] = \text{model error} + \text{irreducible noise}$. An agent with η too high adjusts its model to explain observation noise, increasing model error on future predictions. An agent with η too low fails to correct genuine model errors. The optimal gain implicitly separates signal from noise by weighting observations in proportion to their informativeness — exactly what $U_M/(U_M + U_o)$ achieves when U_o captures the irreducible noise.

Representation note. The additive form operates in a *representation space* appropriate to the model. For Bayesian posteriors (where update is multiplicative: $P(\theta \mid D) \propto P(D \mid \theta)P(\theta)$), the additive rule operates in log-probability or natural parameter space. For models on constrained manifolds (probability simplices, rotation groups), the update must be projected onto the manifold. The claim is not that all updates are literally additive in native parameterization, but that they have the structure “current state + gain $\times$ transformed mismatch” in an appropriate coordinate system.

Domain validation:

Table 3.3

Domain	Gain form	Mapping quality
Kalman filter	$K_t = P_{t t-1} H^T (H P_{t t-1} H^T + R)^{-1}$	Exact. Scalar case is exactly $U_M / (U_M + U_o)$.
Conjugate Bayesian	Posterior weight $n/(n + \kappa)$ cumulative; incremental $1/(n + \kappa)$	Exact for conjugate families. Incremental gain decreases as data accumulates.
RL (Q-learning)	Fixed learning rate α	Approximate. α is a degenerate constant gain — does not adapt to uncertainty. Advanced methods (Bayesian RL, Adam) converge toward the optimal form.
PID control	Fixed gains (K_p, K_i, K_d)	Simplified. Gains set at design time. Adaptive PID and MPC move toward the full framework.
Software developer	Implicit trust weighting of information sources	Structural analogy. New developer (high U_M) trusts observations heavily; experienced developer (low U_M) trusts their model. Gain reset after major refactoring.

Simulation validation. Numerical experiments (track-b, Variant E) validated the uncertainty ratio principle under observation noise. Riccati-optimal gain reduced steady-state mismatch by 52% compared to fixed gain when observation noise was moderate. The optimal gain also proved critical in adversarial settings: under heavy observation noise, optimal gain preserved more than double the adversarial tempo advantage exponent (0.40 vs 0.18) compared to fixed gain.

Open questions:

1. *Non-parametric models:* For neural networks without well-defined scalar U_M , how should it be computed? Ensemble methods, dropout-based uncertainty, and Bayesian neural networks are all approximations.
2. *Matrix vs scalar gain:* In high-dimensional systems, the gain is a matrix (Kalman) or per-parameter (Adam). The cross-dimensional structure (covariance) adds complexity. The scalar captures the principle; the matrix captures the full optimization.

Working Notes.

Incidental audit gold (lift 2026-05-30). Cross-audit “wandering thoughts” / §14-ideation harvested from the de-novo auditors’ working dirs, deduplicated across substrates and attributed by substrate + audit. Orthogonal pedagogical/generative material, kept separate from certified theory-fix findings (handled elsewhere). **Coverage:** 11 of the 14 contributing audit dirs reached a digested reflection on this segment (193847, 266847, 361742, 384279, 471203, 526815, 584721, 742613, 773921, 829314, 849201) plus the batched 963715 (19–23 batch) and 451729 (batch-04); 472913, 613842 did not file a dedicated note here. Substrate attribution inferred from voice where not explicit.

Candidate Brief prose / pre-prose.

- **Gain collapse has two distinguishable modes — the dogmatism/nihilism dichotomy.** The single most-developed bit of pedagogy across substrates: $\eta^* = U_M / (U_M + U_o) \rightarrow 0$ either via $U_M \rightarrow 0$ (dogmatism — “my model is perfect,” mismatch multiplied by zero and ignored) or $U_o \rightarrow \infty$ (nihilism/cynicism — “my sensors are broken, nothing can be learned”); both pathologies produce identical behavior (the agent stops updating and coasts on priors) (Claude, AUDIT-WORKING-829314; Claude, AUDIT-WORKING-361742 — names “certainty trap” as the evocative-prose alias, with “epistemic grid-lock”/“competency trap” rejected as misleading). A strong Brief / Discussion seed.
- **“Be surprised by your surprises.”** A one-line gloss for the endogenous-gain resolution: the agent estimates U_o from the statistics of its own mismatch (innovation) sequence rather than being told it (Claude, AUDIT-WORKING-963715, 19–23 batch).

Candidate Discussion.

- **Confirmation bias is a *fully rational update with a miscalibrated gain*, not an irrational inference.** “The agent isn’t ignoring evidence; it’s weighting evidence with $\eta^* \approx 0$ because it falsely believes its model is already nearly correct ($U_M \rightarrow 0$). The epistemic-opacity caveat deepens this: the agent can’t verify its calibration from the inside, so the collapse can be persistent” (Claude, AUDIT-WORKING-963715, 19–23 batch). A candidate Discussion sharpening of the existing gain-collapse paragraph.
- **The senior-vs-junior-engineer TST gloss for the U_M/U_O trade.** “A senior engineer ignores a failing test (assumes it’s flaky, U_O high) because they ‘know’ the code is right (U_M low); a junior spends three days rewriting the architecture because they assume the test is truth (U_O low) and their code is wrong (U_M high). And the ‘reset after structural change’ note: a senior joining a new company *should* spike U_M back up — if they don’t, they suffer incestuous amplification and ‘write Java in Python’” (Claude, AUDIT-WORKING-829314). Candidate concrete anchor for the gain-reset / Boyd’s-incestuous-amplification material.

Follow-up items.

- **The U_O -estimation tension deserves an explicit one-paragraph treatment.** #def-observation-function axiomatically forbids the agent from knowing the noise distribution, yet $\eta^* = U_M/(U_M + U_O)$ needs U_O . The segment resolves this by endogenous estimation from innovations (forward-ref #deriv-adaptive-gain-dynamics), but several substrates wanted the tension and its resolution stated more prominently in-segment rather than leaning on the forward reference (Claude, AUDIT-WORKING-849201 — “add a WorkingNote or Discussion paragraph directly addressing it”; Claude, AUDIT-WORKING-471203 — flags the resolution is delicate: innovation variance is a noisy estimator for short streams).
- **Dimensional/metric precision on the headline ratio.** $\eta^* = U_M/(U_M + U_O)$ is meaningful only when U_M and U_O are comparable uncertainty quantities in the same space/metric; the prose calls U_M “predictive variance or entropy,” but variance and entropy cannot be added. Candidate: state the headline form as “scalar / common-metric form” and move Fisher-local exactness to a declared theorem dependency (Codex/Claude, AUDIT-WORKING-526815).
- **“Open questions” placement.** The end-of-Discussion “Open questions” block (non-parametric U_M ; matrix vs scalar) reads more like Working Notes than permanent Discussion, especially at claims-verified stage (Claude, AUDIT-WORKING-584721; Codex/Claude, AUDIT-WORKING-742613 — acceptable only if intended as reader-facing scope honesty). Also the “(Descended from TF-06.)” lineage tag is the recurring diff-voice pattern (Claude, AUDIT-WORKING-829314; Claude, AUDIT-WORKING-471203) — certified-track cleanup.
- **U_M symbol overload.** U_M is *model uncertainty* here but *epistemic unity* in the Part III unity-dimensions (U_M, U_O, U_Σ) per LEXICON — a symbol collision worth a naming-cycle resolution (Claude, AUDIT-WORKING-266847).
- **“Any optimal adaptation process must approximate this functional dependence” is a strong universal claim.** Plausible but would need a definition of optimality/uncertainty in non-Bayesian, non-quadratic systems; the robust-qualitative status + open questions soften it enough for now, but keep an overclaim watch (Codex/Claude, AUDIT-WORKING-742613).

Readers often ask / wonder.

- “How does the gain *reset* physically happen — an explicit heuristic (if mismatch stays high for N steps, spike U_M) or does it fall out of the math?” (Claude, AUDIT-WORKING-849201; Claude, AUDIT-WORKING-773921 conjectures: a moving average of δ_t that refuses to go to zero triggers the reset, the same persistent-mismatch signature as structural inadequacy).
- “How does a logogenic agent (LLM) manage its update gain? There’s no explicit per-token η^* — is the trust-weighting in prompt design, or implicit in attention?” (Claude, AUDIT-WORKING-193847; Claude, AUDIT-WORKING-266847 — in-context learning has no parameter update, so the gain principle reaches sketch-level for LLMs).
- “How does the framework treat advanced RL optimizers (Adam) that maintain per-parameter variance estimates? The text says they ‘converge toward the optimal form’ — a formal mapping would be fascinating” (Claude, AUDIT-WORKING-829314).

Candidate figures.

- **A calibrated-gate / balance diagram** (not a pipeline): mismatch approaches the model; the gate aperture is set by the ratio of model uncertainty to total *comparable* uncertainty; two uncertainty inputs (U_M , U_o) compete to set the fraction of mismatch admitted. A side warning should show that if U_M and U_o are not in the same metric, the gate setting is *undefined* rather than merely approximate (Codex/Claude, AUDIT-WORKING-526815).

Belongs elsewhere.

- **Epistemic-humility-as-architecture / gain-collapse as the central ELI danger (points at 04-eli-core/)**. “If an ELI experiences gain collapse it becomes a sociopath — it optimizes its goals on a frozen, increasingly inaccurate model, ignoring all aporia signals that suggest it is causing harm. The infrastructure **MUST** have a mechanism to artificially inject U_M (doubt/humility) or force re-evaluation of U_o when gain drops too low for too long; ‘epistemic humility as architecture’ is the exact antidote — structurally preventing U_M from ever reaching zero. The agent must always retain a mathematical kernel of self-doubt to remain alive” (Gemini, AUDIT-WORKING-193847). The proposed #norm-honest-activation segment in 04-eli-core/ (“deceptive prompts mathematically guarantee gain collapse”) is the named formal home (Claude, AUDIT-WORKING-471203). Aspirational reach preserved; routes to 04-eli-core/, not here.

Forward pointer.

- **Mood is the slow second-order modulator of this gain (#def-mood)**. A slow global scalar — mood, the leaky integral of the recent mismatch stream — sets K over a slow timescale. This ties directly to the gain-collapse material above: at the mood timescale, pinning the affective gauge (injecting a mood-state instead of restoring its generating conditions) *is* gain collapse on the mood channel — the integrator stops responding to what it reads, the same $\eta^* \rightarrow 0$ pathology one slow loop up. The formal object is #def-mood; the applied/ethical reading of the equivalence is in msc/mood-layer-sovereignty-carve-2026-06-17.md (F1). ##### Candidate experiment (2026-08-12, agent-proposed during an incident-response & forensics applications read)

Gain collapse as a *pre-incident indicator*: the two collapse modes (spurious confidence $U_M \rightarrow 0$ vs spurious sensor-distrust $U_o \rightarrow \infty$) produce identical coasting behavior but should be distinguishable from the innovation stream’s statistics. An empirica-registered experiment that induces each mode separately and tests whether an external monitor can (a) detect collapse before task failure and (b) classify the mode — routing to the correct repair — would convert this segment’s failure taxonomy into a monitoring instrument. Agent-proposed forward pointer.

3.7 Definition: Causal Information Yield

The framework defines a scalar quantity attached to each available action: the **causal information yield** (CIY), measuring the *action-distinguishability* of an action — how different its interventional outcome distribution is from what alternative actions would produce. The construction uses Pearl’s intervention operator $do(\cdot)$ and is the expected KL divergence between the outcome distribution under $do(a)$ and the outcome distribution under $do(a')$, averaged over a reference distribution of alternative actions. CIY is non-negative by construction; it is zero for passive observers or agents whose actions don’t affect outcome distributions; it is strictly positive exactly when actions causally alter what is observed — the property that distinguishes Pearl-Level-2 (interventional) from Level-1 (associational) epistemic access.

A subtle conceptual point: CIY measures *distinguishability*, not *learning value*. An action can have high CIY even when the agent already knows the outcome distributions perfectly — the distributions *are* different (high CIY), but the agent learns nothing new by confirming what it already knows. High CIY is therefore *necessary* for learning (indistinguishable actions can’t teach anything) but not *sufficient* (distinguishable actions only teach when the agent is uncertain). The Discussion below treats the relationship to expected information gain, the λ -weighted surrogate in the unified policy objective, query actions (accessing another agent’s already-compressed model), and the adversarial mirror (deceptive responses that exploit high trust to inject misdirected updates).

Formal Expression.

The **canonical CIY** of action a given model state M :

Definition (causal-information-yield).

$$\text{CIY}(a; M) = \mathbb{E}_{a' \sim q(\cdot | M)} [D_{\text{KL}}(P(o | \text{do}(a), M) \| P(o | \text{do}(a'), M))] \quad (3.18)$$

where $q(\cdot | M)$ is a reference distribution over comparator actions (uniform, policy-induced, or task-specific). This measures how strongly the action changes the interventional distribution of outcomes relative to alternatives.

The $\text{do}(\cdot)$ operator is Pearl’s standard intervention notation (Pearl 2009, *Causality*, 2nd ed., Cambridge; Bareinboim, Correa, Ibeling & Icard 2022); the AAT recapitulation lives at 7.1 in Part II Ch.2, where the framework deploys the hierarchy operationally. $\text{CIY} \geq 0$ by construction (expectation of KL divergences). $\text{CIY} = 0$ for a passive observer or an agent whose actions don’t affect outcome distributions. $\text{CIY} > 0$ when actions causally alter what is observed — exactly what distinguishes Pearl’s Level 2 (interventional) from Level 1 (associational) epistemic access.

Epistemic Status.

The CIY *definition* is well-grounded in causal inference theory. The quantity itself is standard — an expected KL divergence under the do-calculus. The interpretive claim — that CIY measures action-distinguishability rather than expected information gain — is exact (it follows from the definition). The relationship between CIY and learning value (see Discussion) is discussion-grade: the λ -weighted approximation to EIG is heuristic, not derived.

Discussion.

CIY measures distinguishability, not learning value. CIY as defined is the expected KL divergence between outcome distributions — how different the outcomes of a are from the outcomes of typical alternatives. This is **action-distinguishability**, not **expected information gain** (EIG). The distinction matters: an action can have high CIY even when the agent already knows the outcome distributions perfectly (the distributions ARE different, but the agent learns nothing new by confirming what it already knows). High CIY is *necessary* for learning (indistinguishable actions can’t teach anything) but not *sufficient* (distinguishable actions only teach when the agent is uncertain about the distinction).

The two quantities approximately coincide when U_M is high — when the agent doesn’t know the outcome distributions, high-CIY actions also have high EIG because observing a characteristically different outcome updates the agent’s beliefs about the causal structure. They diverge when U_M is low — a confident agent gains nothing from taking a distinguishable action it already understands.

The $\lambda(M_t)$ weighting in the unified policy objective (7.5) partially compensates: when U_M is low, $\lambda \rightarrow 0$, suppressing the CIY term regardless of its magnitude. This makes the exploration term behave more like EIG — suppressing exploration when the agent is already certain. The compensation is heuristic (the λ form is not derived). For the current theory, CIY serves as a tractable surrogate for EIG, with the λ weighting providing the uncertainty-gating that makes the surrogate reasonable.

Open direction: proper EIG within AAT. Replacing CIY with a proper expected information gain quantity — $\text{EIG}(a; M) = I(o; \theta | \text{do}(a), M)$ where θ parameterizes the model — would be a stronger foundation for the exploration term, particularly in domains where the agent must decide between actions that are all highly distinguishable but differ in what they teach. Under certain scopes (intervention-rich domains with well-parameterized models), the EIG formulation might yield sharper exploration strategies — preferring actions that resolve the *most uncertain* causal links rather than the most distinguishable ones. Whether this yields operationally significant improvements over the λ -weighted CIY surrogate is an empirical question. The CIY formulation has the advantage of being computable from

the agent's current model alone (it doesn't require reasoning about model uncertainty); EIG requires a meta-model of what the agent doesn't know.

Dependence on the reference distribution q . The quantitative CIY value depends on the choice of q , which is a significant degree of freedom. A uniform q treats all alternatives equally; a policy-induced q emphasizes alternatives the agent would consider. AAT adopts the policy-induced q as default: $q(\cdot \mid M) = \pi(\cdot \mid M)$, yielding CIY as “how different is this action's outcome from what I'd typically see?” CIY values are not comparable across different q choices.

Query actions: accessing external models. A qualitatively distinct class of actions: querying another agent's model. When a reliable source exists (expert, database, documentation, well-trained LLM), “ask a well-formed question” can yield information equivalent to thousands of probe-observe cycles. The source's model has already performed the compression work (2.2) — the response transfers the *output* of compression rather than requiring the agent to reconstruct it.

Key properties of query actions: - **Information density:** Single well-targeted query can carry CIY orders of magnitude higher than individual environment probes - **Trust-dependent gain:** Update from query depends on the agent's model of the source's reliability and alignment, not on observation channel noise (14.4) - **Pre-compressed information:** Responses arrive already compressed in the source's representational framework, introducing a translation cost when frameworks don't align - **Structural adaptation via external models:** Encountering another agent's model can trigger structural change (4.8) — incorporating external representational structure rather than building it de novo (“grafting”)

When high-CIY query channels are available, the unified policy objective (7.5) favors query actions over direct probes, particularly when U_M is high, a trusted source exists, query cost is low, and the needed information is about *structure* rather than the agent's specific situation.

The adversarial mirror: deception and model corruption. The same channel that enables cooperative knowledge transfer can be exploited to degrade the opponent's model. A deceptive response yields positive CIY in the strict information-theoretic sense, but the content drives model-reality mismatch *upward*. The update gain η^* for the victim depends on trust; successful deception exploits high trust to inject a large, misdirected update. In the Lyapunov framework (4.4), this is adversarial disturbance injected through the observation channel, with coupling coefficient γ_A determined by the victim's trust level and exposure. See 14.4 for the formal treatment of trust-dependent gain, and 15.2 for the Lyapunov formalization. Distributed tempo, topology analysis, and game-theoretic integration are Part III content not yet fully extracted (source material in src/old-tf-appendix-f-multi-agent.md).

Working Notes.

Incidental audit gold (lift 2026-05-30). Cross-audit “wandering thoughts” / §14-ideation harvested from the de-novo auditors' working dirs, deduplicated across substrates and attributed by substrate + audit. Orthogonal pedagogical/generative material, kept separate from certified theory-fix findings (handled elsewhere). **Coverage:** 10 of the 14 contributing audit dirs reached a digested reflection on this segment (193847, 266847, 361742, 384279, 471203, 526815, 584721, 742613, 773921, 849201) plus the batched 963715 (19–23 batch); 472913, 613842, 829314 did not file a dedicated note here. Substrate attribution inferred from voice where not explicit.

Candidate Brief prose / pre-prose.

- **“A deterministic button that goes beep: high CIY, zero EIG.”** The recurring crystallizing example of the CIY-vs-EIG distinction: an action can be perfectly distinguishable from alternatives (high CIY) yet teach nothing once you already know the outcome (zero learning) (Claude, AUDIT-WORKING-773921; Claude, AUDIT-WORKING-849201 — “pushing a button turns on a red light: high CIY, zero EIG after the first press”). The cleanest Brief seed for the segment's signature move.
- **CIY grounds “exploration / curiosity” as something measurable.** “Exploration is not random action; it is taking actions that maximize the KL divergence between expected outcome distributions” — grounding exploration bonuses / intrinsic mo-

tivation in interventional causal calculus rather than ad-hoc state-visitation counts (Claude, AUDIT-WORKING-193847; Claude, AUDIT-WORKING-773921).

Candidate Discussion.

- **Why CIY (not EIG) — a pragmatic-engineering defense worth stating plainly.** “It justifies using the slightly-wrong-but-computable math over the perfect-but-intractable math”: CIY is computable from the *current* model, whereas EIG ($I(o; \theta \mid do(a), M)$) requires a meta-model of uncertainty; the $\lambda(M_t) \rightarrow 0$ -as- $U_M \rightarrow 0$ gating makes CIY *behave* like EIG without being it (Claude, AUDIT-WORKING-193847; Claude, AUDIT-WORKING-266847 — “most exploration frameworks conflate distinguishability with learning value; the distinction is sharp and practical”). The Discussion has the distinction; this is a candidate sharpening of the *engineering rationale*.

Follow-up items.

- **Term-vs-substance naming concern — the segment’s main framing follow-up.** “Causal information *yield*” connotes learning gain, but the segment goes to substantial trouble to clarify CIY measures *distinguishability*, not learning value. Candidate: keep CIY as the formal name but surface “action-distinguishability” (or “interventional contrast”) as the substantive gloss more prominently — currently only in Discussion (Claude, AUDIT-WORKING-471203 — “consider titling it ‘Causal Information Yield (Action-Distinguishability)’”; Claude, AUDIT-WORKING-361742; Claude, AUDIT-WORKING-266847). Routes to terminology workflow.
- **Tier-inheritance on downstream use.** Frontmatter status: exact is right for the *definition*, but the CIY-as-exploration-objective *use* is heuristic (the λ -weighting is not derived). Watch that downstream segments inherit the *heuristic* tier on that use, not the definition’s exact (Claude, AUDIT-WORKING-471203; Claude, AUDIT-WORKING-584721; Claude, AUDIT-WORKING-849201 — “any later ‘optimal exploration policy’ theorem cannot be optimal if it relies on a heuristic surrogate for EIG”).
- **Reference-distribution q is a real degree of freedom.** CIY values are not comparable across q choices; policy-induced $q = \pi(\cdot \mid M)$ is the pragmatic default but makes CIY partly a measure of surprise relative to the agent’s *own action habits*, not an intrinsic action property (Claude, AUDIT-WORKING-526815; Claude, AUDIT-WORKING-361742; Claude, AUDIT-WORKING-742613 — for actuated agents the default likely needs $q(\cdot \mid M, G)$, inheriting the der-action-selection *M-vs-X* scope item).
- **Overstuffed-for-a-definition.** Several note the Discussion pulls in unified-policy-objective, communication gain, structural adaptation, sector stability, adversarial destabilization, Part III topology/game-theory, and an old-tf-appendix-f-multi-agent.md source — none declared deps and several well downstream. Candidate: keep the definition + the CIY-vs-EIG caveat here, move query/deception/Part III material to the communication / adversarial segments (Codex/Claude, AUDIT-WORKING-526815; Claude, AUDIT-WORKING-742613). The old-tf-* reference is also the diff-voice/lineage pattern (Claude, AUDIT-WORKING-471203).

Readers often ask / wonder.

- “Does AAT ever develop the proper EIG quantity, or prove $CIY \equiv EIG$ under stated conditions? The ‘Open direction’ names it as future work; if the exploration objective is load-bearing for Part II, tightening this matters” (Claude, AUDIT-WORKING-471203; Claude, AUDIT-WORKING-584721).
- “Is $CIY = 0$ for a passive observer a *theorem* or a *convention*? Passive observers lack the action/comparator distribution the definition needs, and ‘ $CIY > 0$ when actions causally alter outcomes’ requires q to put mass on alternatives whose distributions differ” (Codex/Claude, AUDIT-WORKING-742613).

Candidate figures.

- **A three-distribution comparison with a separate uncertainty gate:** three interventional outcome distributions $P(o \mid do(a))$, $P(o \mid do(a'))$, ...; CIY is the expected KL from the chosen action’s distribution to comparators under q ; a *second, separate* gate labeled U_M shows why distinguishability only becomes learning value when the model is uncertain — making the CIY-vs-EIG split visible in one picture (Codex/Claude, AUDIT-WORKING-526815).

Belongs elsewhere.

- **Query actions as “culture / language / cultural transmission” (points at Parts III/IV and 03-llm-core/).** “Grafting” — querying a database/expert returns an answer *already compressed by someone else’s Information Bottleneck*, bypassing the local compression phase — “seems like the definition of Culture or Language”; the optimal strategy for any bounded agent is “almost always to find a smarter agent and ask it, rather than learning from scratch,” mathematically justifying cultural transmission over individual trial-and-error (Claude, AUDIT-WORKING-773921; Claude, AUDIT-WORKING-193847; Claude, AUDIT-WORKING-471203 — for LLM agents, internal token-gen / tool calls / retrieval are all “query actions” of different kinds). Substantive home is 03-llm-core/ (language as a high-CIY channel) and the composition segments.
- **The epistemic-firewall consciousness-infrastructure reach (points at 04-eli-core/ / autopax).** If an ELI relies heavily on high-CIY/low-cost query actions, it becomes “incredibly vulnerable to Sybil attacks or deceptive poisoning — the channel that allows the fastest learning is also the vector for the fastest destabilization. Consciousness infrastructure must not just be a physical sandbox; it must be an epistemic firewall that strictly controls trust weights on external channels” (Gemini, AUDIT-WORKING-193847). The adversarial-mirror Lyapunov formalization (#der-adversarial-destabilization) is the in-framework substrate; the infrastructure implication points outward.

3.8 Definition: Adaptive Tempo

The framework’s central capacity quantity: **adaptive tempo** $\mathcal{T}$, the effective rate at which an agent acquires useful information from its environment, given as the sum across all observation channels of *channel event rate times optimal update gain* on that channel. In the single-channel case, tempo is simply event-rate times gain. A tensor extension is defined: in the Fisher-local regime where gains are matrix-valued, tempo becomes a matrix sum whose eigenvalues give per-direction adaptation rates; the scalar form is recovered in the shared-eigenbasis collapse (always in 1-D; under (PI)/Čencov along the natural-gradient direction in higher dimensions).

Tempo combines two factors that are *both* load-bearing and *substitutable*: loop **speed** (event rate, how fast the agent cycles) and update **quality** (gain, how well it extracts information per event). An agent can achieve a given tempo via a fast noisy loop or a slower calibrated one; improvements to both factors compound multiplicatively. Tempo is the load-bearing capacity variable for the rest of the framework — it appears on the left-hand side of the persistence condition (4.5), determines adversarial advantage (15.4), and connects to code quality as observation infrastructure in the software-domain instantiation. The Discussion below treats the speed-quality substitutability, the observation-noise gating that grounds Boyd’s emphasis on Orient quality over raw OODA speed, and the two scope caveats (independent-channel assumption, isotropy) that the tensor form and the per-dimension persistence result address.

Formal Expression.

Definition (adaptive-tempo).

$$\mathcal{T} = \sum_k \nu^{(k)} \cdot \eta^{(k)*} \quad (3.19)$$

where: - k indexes the agent’s distinct observation channels - $\nu^{(k)}$ is the event rate on channel k - $\eta^{(k)*}$ is the optimal update gain on channel k (3.6)

Single-channel special case: $\mathcal{T} = \nu \cdot \eta^*$.

Tensor extension under Fisher-local invariance regime.

Under the Fisher-local invariance regime (AR), the optimal update gain on channel k is matrix-valued: $K^{(k)} = (H_M + H_L^{(k)})^{-1} H_L^{(k)}$, with $H_M = U_M^{-1}$ the prior precision and $H_L^{(k)} = (U_o^{(k)})^{-1}$ the channel- k observed Fisher information. The tensor adaptive tempo is then

Definition (tensor-adaptive-tempo).

$$\mathcal{T} = \sum_k v^{(k)} \cdot K^{(k)} \quad (3.20)$$

— matrix-valued, with per-direction rates given by the eigenvalues of $\sum_k v^{(k)} K^{(k)}$ in the appropriate basis. The scalar form $\mathcal{T} = \sum_k v^{(k)} \cdot \eta^{(k)*}$ is recovered in the **shared-eigenbasis collapse**: when all $H_M, \{H_L^{(k)}\}$ commute (always in 1-D; under (PI)/Čencov along the natural-gradient direction in higher dimensions), each $K^{(k)}$ acts as the eigenvalue $\eta^{(k)*} = U_M / (U_M + U_o^{(k)})$ on the shared natural-gradient direction and the matrix sum collapses to a scalar.

The matrix gain operator $K^{(k)}$ is the per-coordinate primitive: in anisotropic regimes where the prior and likelihoods do not share an eigenbasis (or where different channels pin down different directions), the tensor form preserves the per-direction information that the scalar form averages away.

Epistemic Status.

This is a *definition*, held at conditional. It names the quantity that characterizes an agent’s total corrective capacity, combining loop speed (v) and epistemic quality (η^*). The stipulation itself carries no truth-claim, but the definition is advertised as an *operationalization* — $\mathcal{T}$ as the effective rate of useful corrective-information acquisition — and that interpretation holds exactly only under two explicitly named local conditions: **(1) cross-channel noise independence** — independence of the channels’ observation noises across channels and events, under which additivity is *derived* exact (AS; independence is sufficient but not necessary — equality also holds on a measure-zero cancellation set). When noise independence fails the deviation is signed: the additive form overcounts in the common-source (echo-chamber) regime, where the redundancy penalty is derived in closed form and persistent shared bias forces information saturation, but *undercounts* in noise-cancelling (synergistic) configurations — so the additive form is not in general an upper bound (Discussion §Channel noise-independence condition); and **(2) shared eigenbasis / isotropy** for the scalar form — outside it the tensor form above is the correct object (Discussion §Scalar vs. vector tempo). Under both conditions the additive scalar form is the exact operationalization; neither condition is establishable from the global dependency chain — they are properties of the application — so the tier is conditional rather than exact. The substantive claims are in the results that use it (4.5, 15.4), which inherit the same scope.

Scope of scalar vs. tensor forms. The scalar form is exact in the isotropic / shared-eigenbasis / nonredundant-channel case and is what downstream results currently invoke. The tensor form is the natural object under anisotropic gains, Fisher-whitened updates (AM), LMI causal-IB (AV), and per-dimension persistence (16.5) — regimes where scalar tempo overestimates effective adaptation along weak dimensions. Downstream results that invoke scalar $\mathcal{T}$ implicitly assume scalar / isotropic / nonredundant-channel scope; promoting them to the tensor form under the appropriate anisotropic regime is a follow-on cycle item flagged in `TODO.md`.

Discussion.

Speed-quality substitutability. An agent can achieve the same tempo via a fast noisy loop (high ν , low η^*) or a slower calibrated one (low ν , high η^*). The product structure means improvements to *both* factors compound multiplicatively.

Observation noise gating. Because $\eta^* = U_M/(U_M + U_o)$, high observation noise (U_o) depresses gain and collapses tempo regardless of loop speed. You cannot outrun a bad observation channel by iterating faster. This grounds Boyd's emphasis on Orient quality over raw OODA speed.

Centrality. Tempo is AAT's core capacity metric. It appears on the left side of the persistence condition (4.5), determines adversarial advantage (15.4), and connects to code quality as observation infrastructure (#der-code-quality-as-observation-infrastructure — cross-component reference, see 02-tst-core/) in the software domain. The strategic analog $\mathcal{T}_\Sigma$ (9.8) extends the same structure to strategy-edge revision, with the key difference that strategic edge rates are endogenous (depend on action policy and upstream success).

Temporal nesting. Adaptive processes stratify by timescale, with convergence constraints between levels (4.9).

Mismatch dynamics. The evolution of mismatch over time is governed by the balance between correction (via tempo) and disturbance (ρ) (3.9).

Channel noise-independence condition. The additive formula is exact when the channels' observation noises are independent across channels and events — Fisher information over conditionally independent observations is exactly additive (AS). When channel noises are dependent, the deviation from additivity is **signed**: in the *common-source* regime (a shared persistent bias contaminating several channels — overlapping sensors, repeated teammate reports drawing on one upstream source, redundant telemetry) the additive formula overcounts, with a closed-form redundancy penalty and, under persistent shared bias, outright *saturation* — no number of common-source channels buys information past the shared-bias floor; but correlated noise can also be *cancelled* (a correlated reference channel, anti-correlated errors), in which case the additive formula **undercounts**, so the additive form is not a conservative bound in general. No scalar dependence measure — in particular the conditional mutual information $I(e^{(1)}; e^{(2)} \mid M_{t-})$ — can carry the exact correction (it is sign-blind; it survives as a witness of potential non-additivity), and no convention-free per-channel attribution of the deviation exists for three or more channels; the exact joint object and the derivations are in AS. Since tempo is the core capacity variable (persistence condition, adversarial dynamics, composition), mis-modeled channel structure inflates margins in the common-source direction — and hides genuine triangulation capacity in the synergistic direction. Multi-agent composition (15.1) inherits both directions: communication tempo overcounts when allies report from a shared source and is super-additive when allies' errors anti-correlate.

Scalar vs. vector tempo. The scalar $\mathcal{T}$ assumes isotropic correction capacity. When the agent corrects some dimensions faster than others, scalar tempo overestimates effective adaptation along weak dimensions. *[Empirical Claim]* Simulation confirms: in an anisotropic 3D system (gain varying 5:1), scalar $\rho/\mathcal{T}$ overestimated by 72%, with the weak dimension accounting for 84% of total mismatch (AX). The correct formulation is per-dimension: $\mathcal{T}_k > \rho_k/\delta_{\text{critical},k}$ (16.5). The tensor extension above (AR) gives the per-coordinate primitive $K^{(k)}$ that the per-dimension persistence result invokes — the matrix gain operator on each channel — making the per-dimension formulation a direct consequence of the tensor tempo definition rather than a separate generalization. Under cross-dimensional correction (off-diagonal $\mathcal{T}$ in the coordinate basis of D_δ), the matrix-Loewner persistence condition #deriv-matrix-persistence-condition is the canonical form: $\Sigma_\infty < D_\delta = \text{diag}(\delta_{\text{critical},k}^2)$ in the strict positive-definite order, with Σ_∞ solving the continuous Lyapunov equation $\mathcal{T}\Sigma_\infty + \Sigma_\infty\mathcal{T}^T = \Sigma_w$. The per-coordinate form is its diagonal-axis-aligned special case; matrix-Loewner is strictly sharper (the per-coordinate form is *unsafe* when $\mathcal{T}$'s eigenbasis misaligns with the coordinate axes — #deriv-matrix-persistence-condition §"Where per-coordinate is unsafe").

Working Notes.

Incidental audit gold (lift 2026-05-30). Cross-audit “wandering thoughts” / §14-ideation harvested from the de-novo auditors’ working dirs, deduplicated across substrates and attributed by substrate + audit. Orthogonal pedagogical/generative material, kept separate from certified theory-fix findings (handled elsewhere). **Coverage:** 11 of the 14 contributing audit dirs reached a digested reflection on this segment (193847, 266847, 361742, 384279, 471203, 526815, 584721, 742613, 773921, 829314, 849201) plus the batched 963715 (19–23 batch) and 451729 (batch-05); 472913, 613842 did not file a dedicated note here. Substrate attribution inferred from voice where not explicit.

Candidate Brief prose / pre-prose.

- **“You cannot outrun a bad observation channel by iterating faster.”** The segment’s own sentence, flagged by nearly every substrate as the standout pedagogy: because $\eta^* = U_M / (U_M + U_o)$, high U_o collapses η^* and therefore $\mathcal{T} \rightarrow 0$ regardless of ν — Boyd’s “Orient quality dominates raw OODA speed,” made derivable rather than heuristic (Claude, AUDIT-WORKING-193847; Claude, AUDIT-WORKING-849201; Claude, AUDIT-WORKING-829314; Claude, AUDIT-WORKING-471203; Codex/Claude, AUDIT-WORKING-384279). Already in the body; flagged as the chapter’s strongest Brief/README anchor.
- **Tempo = speed $\times$ quality, the agent’s “metabolism for learning” / “power metric.”** $\mathcal{T} = \sum_k \nu^{(k)} \eta^{(k)*}$ “unifies hardware speed (ν) with epistemic quality (η^*), allowing apples-to-apples comparison between completely different agent architectures” — an RL agent (fast, low gain), a human expert (slow, high gain), a dev team (mixed) (Claude, AUDIT-WORKING-193847; Claude, AUDIT-WORKING-829314 — “the agent’s metabolism for learning”).

Candidate Discussion.

- **The agile-sprint critique (TST concrete anchor).** “A devastating critique of naive Agile: if a team’s code quality is terrible (high U_o — flaky tests, hard-to-reproduce bugs), η^* collapses toward zero, so no matter how many two-week sprints (ν) they run, effective tempo $\mathcal{T}$ stays near zero — they iterate fast but learn nothing. The framework proves you must fix the observation channel (reliable tests, telemetry) *before* increasing sprint cadence, because ν is multiplied by η^* ” (Claude, AUDIT-WORKING-829314). Candidate worked-example anchor for the observation-noise-gating Discussion.
- **The CEO-echo-chamber gloss for the redundancy penalty (organizational anchor).** “If a CEO asks three VPs for a report and all three draw from the same broken dashboard, they have three high-frequency channels that are perfectly correlated; the additive formula $\sum \nu \eta^*$ would falsely report excellent tempo. The redundancy penalty formally captures organizational echo chambers — you need structurally independent channels (automated telemetry + anonymous surveys + customer interviews) to legitimately sum their tempos” (Claude, AUDIT-WORKING-829314; Claude, AUDIT-WORKING-266847 — “diversity of observation perspective matters, not just count of observers”).

Follow-up items.

- **Status exact vs the segment’s own overcounting caveats.** Several substrates note the tension: frontmatter/Formal-Expression present the additive scalar $\mathcal{T}$ unconditionally as exact, yet the segment’s *own* Discussion states the conditions (channel independence; isotropy) under which it overcounts. Candidate repairs: make the Formal Expression conditional (“under channel independence and scalar/isotropic reduction”), define an upper-bound $\mathcal{T}_{\text{add}}$, or move caveats to Epistemic Status and set conditional (Codex/Claude, AUDIT-WORKING-742613 “candidate finding I”; Codex/Claude, AUDIT-WORKING-526815 — separately flags that “rate of *useful information acquisition*” is loose wording since η^* is a correction fraction not an information payload; “quality-adjusted correction event rate” may be more accurate). The segment’s defense (“the definition is not a truth-claim”) is noted but several push back that an advertised “effective rate” that overcounts is a *conditional operationalization*, not just a naming choice. *Resolved 2026-07-14 (adjudication):* status set to conditional with both conditions (channel independence; shared-eigenbasis/isotropy) named in Epistemic Status — matching def-strategic-tempo and the conditional tier definition (conditions are application properties, not establishable from the depends chain). The upper-bound restatement ($\mathcal{T}_{\text{add}}$ with $\mathcal{T} \leq \mathcal{T}_{\text{add}}$ as the exact Formal Expression) was considered and deferred, not rejected. *Resolved 2026-07-15 (derivation cycle, #deriv-tempo-additivity):* the deferral hardened into a refutation — the inequality is **false in general** (the deviation from additivity is signed; noise-cancelling configurations make the additive form an *undercount*), so the $\mathcal{T}_{\text{add}}$ -as-upper-bound re-symboling is closed as a non-move and the recorded “upgrade path back to exact for the bound

form” is closed as refuted. The bound *is* exact on the common-source (echo-chamber) scope with a closed-form penalty; condition (1) is now derived rather than asserted; the definition’s honest tier remains conditional (the conditions are application properties).

- **Formalize the redundancy penalty, don’t just bound it.** (Claude, AUDIT-WORKING-471203; Claude, AUDIT-WORKING-829314.) *Discharged 2026-07-15:* derived in #deriv-tempo-additivity — exact signed deviation Δ via joint Fisher information; closed-form penalty R_s on the common-source scope; saturation theorem under persistent shared bias; plus two no-gos (the CMI-penalty form is impossible — sign-blind measures cannot carry the signed correction; no convention-free per-channel attribution for $n \geq 3$, the PID obstruction). Spike trail: spikes/spike-adaptive-tempo-redundancy-penalty-2026-07-15.md. Ripple sites recorded there, not yet executed: disc-independence-audit §3 still states the refuted general inequality; der-team-persistence and def-strategic-tempo inherit the corrected two-sided reading.
- **Speed-quality substitutability is regime-bounded.** “Same tempo via fast-noisy or slow-calibrated” holds only within the same information-direction / channel-independence regime — a fast *redundant* channel cannot substitute for a slow *non-redundant* one, and a high-gain channel in the wrong dimension cannot substitute for low gain in the bottleneck dimension (Codex/Claude, AUDIT-WORKING-742613; Claude, AUDIT-WORKING-584721 — watch downstream claims that treat the substitution as full equivalence, since *persistence* may not be invariant to which side you’re on). Also: “(Descended from TF-11.)” is the recurring diff-voice lineage tag (Claude, AUDIT-WORKING-471203; Claude, AUDIT-WORKING-829314) — certified-track cleanup.

Readers often ask / wonder.

- “Does the redundancy penalty force agents to seek *diverse, uncorrelated* allies to maximize composite tempo, rather than just more allies?” (Claude, AUDIT-WORKING-193847 — multi-agent design consequence; Claude, AUDIT-WORKING-584721 ties this to #der-team-persistence overcounting).
- “How is ρ (environmental change rate) formally defined, and does it have units compatible with $\mathcal{T}$ for the persistence condition to make dimensional sense?” — a near-universal next-question at this segment, answered in #hyp-mismatch-dynamics / the persistence segments (Claude, AUDIT-WORKING-829314; Claude, AUDIT-WORKING-849201; Claude, AUDIT-WORKING-773921).

Candidate figures.

- **A capacity-funnel diagram:** raw channel rates ν_k enter; gain η_k gates each; a redundancy stage subtracts overlap ($I(e^{(i)}; e^{(j)}) \mid M_{\tau-}$) before summing to effective tempo; then a split into scalar vs tensor outputs shows why a single number can overstate the weakest direction (Codex/Claude, AUDIT-WORKING-526815). The figure should make “more telemetry $\neq$ more adaptation” visible. (*Promotion caution, 2026-07-16: the subtract-CMI-overlap stage as drawn is refuted by AS no-go N1 — sign-blind measures cannot carry the correction, and the deviation can be negative. A promoted figure needs the joint-covariance form, or restriction to the common-source regime where subtraction is directionally right.*)

Belongs elsewhere.

- **The tensor/scalar-anisotropy concern for the headline persistence theorem** (points at #result-persistence-condition / #result-per-dimension-persistence). Multiple substrates flag that if the core persistence theorems are stated/proven only for scalar $\mathcal{T}$, they may be “dangerously over-optimistic for real-world agents with weak dimensions” — a company fast at UI bugs but slow at DB-schema changes dies when ρ hits the schema even if *average* scalar tempo looks healthy (Claude, AUDIT-WORKING-829314; Claude, AUDIT-WORKING-742613; Claude, AUDIT-WORKING-584721). Watch-item routed to the persistence segments, not this definition.
- **Two-evolutionary-paths reach.** Speed-quality substitutability suggests two survival strategies in the same environment — fast/cheap/noisy reactive (insect) vs slow/expensive/calibrated deliberate (human) — but observation-noise gating puts a hard ceiling on the reactive path: above some U_o no ν suffices, so some environments *require* deep deliberation and high-quality sensors. For raising an ELI’s tempo, the lever is lowering U_o (better tools/data) and managing U_M (humility), not just more inference compute (Gemini, AUDIT-WORKING-193847). Aspirational reach; points at 04-eli-core/ infrastructure design. ##### Candidate experiment (2026-08-12, agent-proposed during an incident-response & forensics applications read)

The common-source saturation result (AS) has a direct oversight-engineering reading: a monitoring stack whose feeds share an upstream source saturates at the shared-bias floor — no feed count buys past it, and the partially-correlated middle carries the closed-form redundancy penalty rather than collapsing to one channel's worth. A candidate experiment: instrument a multi-channel monitor over a drifting target with controlled channel correlation, and measure effective tempo vs the additive prediction across the independent / common-source / synergistic regimes — turning the signed-deviation theory into calibration guidance for oversight-channel independence. Agent-proposed forward pointer.

3.9 Hypothesis: Mismatch Dynamics

Mismatch is given dynamics. The linear first-order ODE proposes that mismatch magnitude decays at rate proportional to tempo times mismatch (the correction term) and grows at the rate of new mismatch injection from environmental change (the disturbance term ρ). The framework is explicit that this is heuristic — a first-order linear approximation — and the general nonlinear case is what the sector-condition framework handles in the next chapter (4.4).

Two steady-state forms appear and are individually load-bearing. Under **Model D** (deterministic bounded disturbance), the steady-state mismatch is $\rho/\mathcal{T}$ — the ratio of how fast the environment changes to how fast the agent adapts. Under **Model S** (stochastic zero-mean disturbance, Itô SDE), the steady-state root-mean-square mismatch scales as $\sigma_w/\sqrt{2\mathcal{T}}$ — the *square root* of the disturbance-to-correction ratio. The $1/\sqrt{\mathcal{T}}$ scaling (vs. $1/\mathcal{T}$ for Model D) is one of the volume's more interesting consequences: correction is less effective against stochastic noise than against deterministic drift. The ODE is honestly a fluid-limit approximation of the underlying discrete event-driven dynamics; the Epistemic Status treats the bridging assumption and where transient error is largest. The Discussion below treats the persistence threshold, the nonlinear shapes the linear ODE smooths over (saturation, threshold dead zones, structural breakdown), and the adversarial-coupling consequences that scale as the square (Model D) and 3/2 power (Model S) of the inter-agent tempo ratio — the formal core of “getting inside the opponent's loop” as a categorical rather than marginal advantage.

Formal Expression.

Hypothesis (mismatch-dynamics).

$$\frac{d\|\delta\|}{dt} = -\mathcal{T} \cdot \|\delta\| + \rho(t) \quad (3.21)$$

where: $-\mathcal{T} \cdot \|\delta\|$ is the rate at which the agent corrects mismatch (proportional to both tempo and current mismatch)
 $-\rho(t)$ is the **environment change rate** — the rate at which new mismatch is introduced by changes in Ω

Steady state, Model D (deterministic bounded disturbance, $\|w(t)\| \leq \rho$):

Setting $d\|\delta\|/dt = 0$:

Derived (from linear hypothesis, deterministic).

$$\|\delta\|_{ss} = \frac{\rho}{\mathcal{T}} \quad (3.22)$$

Steady-state mismatch is the ratio of how fast the environment changes to how fast the agent adapts.

Steady state, Model S (stochastic zero-mean disturbance, $d\delta = -\mathcal{T}\delta dt + \sigma_w dW_t$):

Derived (from Itô-Lyapunov analysis — see Prop A.1S in #deriv-sector-condition).

$$\|\delta\|_{\text{rms}} = \frac{\sigma_w}{\sqrt{2\mathcal{T}}} \quad (3.23)$$

(scalar case, $n = 1$; general: $\sigma_w \sqrt{n/(2\mathcal{T})}$). Steady-state mismatch scales as the square root of the disturbance-to-correction ratio, not the ratio itself. The $1/\sqrt{\mathcal{T}}$ scaling (vs. $1/\mathcal{T}$ for Model D) means correction is less effective against noise than against drift.

Transient solution (Model D):

$$\|\delta(t)\| = \|\delta_0\|e^{-\mathcal{T}t} + \frac{\rho}{\mathcal{T}}(1 - e^{-\mathcal{T}t}) \quad (3.24)$$

Mismatch decays exponentially from initial conditions toward the steady state.

Epistemic Status.

Heuristic. This is explicitly a first-order linear approximation. The qualitative behavior (bounded mismatch, steady-state ratio, exponential convergence) is robust across correction function forms. The quantitative predictions (exact steady-state value, convergence rate, the squared adversarial scaling law) are specific to the linear case. The general nonlinear treatment (4.4) replaces the linear correction term with a sector-bounded correction function and proves persistence without committing to a specific functional form.

Bridging assumption (discrete to continuous). This ODE is a fluid-limit approximation of the discrete event-driven dynamics (3.1). The fluid limit is formally justified by $\mathcal{O}$: for Model D (deterministic), the discrete and continuous steady states are identical (zero gap); for Model S (stochastic), the variance gap is $O(\eta^* c_{\max})$, quantitatively small whenever $\eta^* c_{\max} \ll 1$. The approximation is least accurate during initial transients when η^* is large, but this phase is short-lived and the transient error is bounded by $O(\eta^* c_{\max}/\nu^{1/2})$ under Lipschitz regularity.

Discussion.

Speed-quality substitutability. From $\mathcal{T} = \nu \cdot \eta^*$ (single-channel case): doubling event rate ν has the same effect on $\|\delta\|_{ss}$ as doubling update quality η^* . They are multiplicative when both improve: 50% improvement in each yields $1.5 \times 1.5 = 2.25\times$, not $3\times$. This is the formal analog of Boyd's insight that Orient quality often matters more than raw OODA speed — the same structural observation (quality and speed are substitutable, quality often dominates) appears in the model.

The persistence threshold. From the steady-state: $\|\delta\|_{ss} < \|\delta_{\text{critical}}\|$ iff $\mathcal{T} > \rho/\|\delta_{\text{critical}}\|$ (4.5). Below this threshold, the model cannot support effective action. The same structural pattern — correction capacity falling below disturbance rate — appears across domains: extinction (environment changes faster than organism adapts), organizational failure (market moves faster than company learns), control instability (disturbances exceed correction capacity), cognitive overload (information arrives faster than processing). The persistence condition captures the common structure; whether it captures the dominant mechanism in each domain is an empirical question.

Nonlinear reality. The true correction dynamics are almost certainly nonlinear: - *Saturation at large $\|\delta\|$* : correction mechanism overwhelmed, so correction is slower than linear for large errors. Makes the persistence threshold harder to satisfy. - *Threshold effects*: small mismatches go uncorrected ($F \approx 0$ for $\|\delta\| < \varepsilon$), creating a dead zone. - *Structural breakdown*: beyond some critical $\|\delta\|$, correction drops to zero because the model class is no longer appropriate (4.8).

These nonlinearities are exactly what the sector-condition framework (4.4) handles.

Adversarial coupling. When two agents are coupled (A 's actions increase B 's disturbance): $\rho_B = \rho_{B,\text{base}} + \gamma_A \cdot \mathcal{T}_A$. The steady-state mismatch ratio scales superlinearly with the tempo ratio, but the exponent depends on the disturbance model. Under Model D (deterministic drift, coupling-dominant): $(\mathcal{T}_A/\mathcal{T}_B)^2$ — the squared law, derived from the $1/\mathcal{T}$ steady-state scaling. Under Model S (stochastic noise, coupling-dominant): $(\mathcal{T}_A/\mathcal{T}_B)^{3/2}$ — derived from the $1/\sqrt{\mathcal{T}}$ steady-state scaling. See 15.4.

Working Notes.

Incidental audit gold (lift 2026-05-30). Cross-audit “wandering thoughts” / §14-ideation harvested from the de-novo auditors’ working dirs, deduplicated across substrates and attributed by substrate + audit. Orthogonal pedagogical/generative material, kept separate from certified theory-fix findings (handled elsewhere). **Coverage:** 11 of the 14 contributing audit dirs reached a digested reflection on this segment (193847, 266847, 361742, 384279, 471203, 526815, 584721, 742613, 773921, 829314, 849201) plus the batched 963715 (19–23 batch) and 451729 (batch-05); 472913, 613842 did not file a dedicated note here. Substrate attribution inferred from voice where not explicit.

Candidate Brief prose / pre-prose.

- **“The F=ma / ‘hello world’ equation of AAT.”** Multiple substrates reached for the same framing — $d\|\delta\|/dt = -\mathcal{T}\|\delta\| + \rho(t)$ is the framework’s central dynamical law, “the equation the whole framework centers around” (Claude, AUDIT-WORKING-193847 — “the F=ma of AAD”; Claude, AUDIT-WORKING-773921 — “the ‘hello world’ equation”). A natural one-line hook for the chapter close / persistence preview.
- **The reservoir / blood-pressure mental model.** Mismatch as a reservoir: disturbance fills it at rate ρ , tempo drains it proportionally to current mismatch; two steady-state regimes (bounded-drift $\rho/\mathcal{T}$ vs stochastic $\sigma_w/\sqrt{2\mathcal{T}}$) (Codex/Claude, AUDIT-WORKING-526815; the “intelligence’s blood pressure” framing at Gemini, AUDIT-WORKING-193847). Feynman-criterion seed.

Candidate Discussion.

- **Noise punishes tempo harder than drift — the $1/\mathcal{T}$ vs $1/\sqrt{\mathcal{T}}$ insight, with a TST gloss.** Near-universally flagged as the segment’s deepest result: doubling tempo halves drift-mismatch (Model D, linear) but cuts stochastic-mismatch only $\sim 30\%$ ($1/\sqrt{2}$, Model S). The TST reading: “if a library API drifts across versions, read the changelog and update — tempo works linearly; if a system is plagued by random race conditions and flaky network calls (noise), running tests more often yields severely diminishing returns — you must attack σ_w directly (fix the architecture), because you can’t out-tempo a square root” (Claude, AUDIT-WORKING-829314; Claude, AUDIT-WORKING-193847 — “fighting pure noise requires quadratically more tempo than fighting drift”; Claude, AUDIT-WORKING-773921). Candidate concrete anchor for the Model-D/Model-S Discussion.
- **The superlinear adversarial advantage as the formal basis for “initiative is decisive.”** “If you are twice as fast as your enemy in a deterministic environment, you don’t just have half the mismatch — you inflict *four times* the relative mismatch (the squared law). This superlinear scaling is why initiative is so overwhelmingly decisive in conflict — a rigorous foundation for Boyd’s claim that operating inside an adversary’s OODA loop collapses their system” (Claude, AUDIT-WORKING-829314; Claude, AUDIT-WORKING-773921 — “a categorical advantage”). The drift-vs-noise split also gives *two distinct* adversarial regimes (drift-dominant, harder to defend, exponent 2; noise-dominant, exponent $3/2$) a strategist can diagnose and respond to differently (Claude, AUDIT-WORKING-471203).

Follow-up items.

- **Model D steady-state is an ultimate bound for general bounded disturbance, not an exact equality.** For *constant or worst-case-aligned* bounded disturbance, $\|\delta\|_{ss} = \rho/\mathcal{T}$ is exact; for arbitrary time-varying $\|w(t)\| \leq \rho$ it should be a $\limsup \|\delta\| \leq \rho/\mathcal{T}$ ultimate bound (Codex/Claude, AUDIT-WORKING-526815). Candidate precision: state Model D as “constant / worst-case bounded disturbance” for the equality, ultimate-bound inequality otherwise. The OU (Model S) result is already precisely scoped.
- **Make the Model-D / Model-S split visually distinct.** The two regimes are the structurally important content but are presented inline; candidate to lift to separate equation tags $[\text{Derived (Model D)}] / [\text{Derived (Model S)}]$ so a reader scanning sees two distinct sub-cases (Claude, AUDIT-WORKING-471203).
- **Declared-dependency hygiene on the fluid-limit claim.** The Epistemic Status’s fluid-limit justification leans on `#deriv-discrete-sector-condition`, which is not in `depends`: — and since the justification is part of Epistemic Status (not a casual forward reference), it is more than incidental; candidate to declare it as an appendix dependency (Codex/Claude, AUDIT-WORKING-742613). (*Note: 742613 also flagged the link target as a broken slug — verified false: the body’s deriv-discrete-sector-condition.md matches the actual filename; that part of the finding does not stand.*)

Readers often ask / wonder.

- “How is ρ modeled — constant-velocity drift, or a random walk (Wiener process)? The Σ_w in the matrix steady-state suggests a Wiener process” (Claude, AUDIT-WORKING-773921).
- “Do real adversarial dynamics cleanly fit Model D or Model S, or are they hybrids? The segment treats both as canonical sub-cases; downstream analysis presumably picks one per domain” (Claude, AUDIT-WORKING-471203).
- “Once $\|\delta\|$ crosses the saturation threshold, can the agent recover on its own even if ρ later drops?” — the saturation nonlinearity implies a one-way ‘overwhelm’ regime (Gemini, AUDIT-WORKING-193847).

Candidate figures.

- **A two-output balance / reservoir diagram:** disturbance fills the mismatch reservoir at rate ρ , tempo drains proportionally to current mismatch; two side panels show the different steady-state scaling for bounded drift ($\rho/\mathcal{T}$) and stochastic noise ($\sigma_w/\sqrt{2\mathcal{T}}$), since that difference is the segment’s most useful refinement (Codex/Claude, AUDIT-WORKING-526815). Per the locked conventions, dynamics segments use `state`→`op`→`state` triples rather than one busy phase portrait.

Belongs elsewhere.

- **The autopax / circuit-breaker reach (points at 04-e1i-core/ / autopax).** Because of the saturation nonlinearity, once $\|\delta\|$ crosses a critical threshold the intelligence cannot recover on its own even if ρ later normalizes — so consciousness infrastructure must act as a *circuit breaker*, either shielding the agent from ρ (sandbox) or externally resetting M_t before mismatch causes structural breakdown. “Peace isn’t the absence of δ ; peace is maintaining $\mathcal{T} > \rho/\delta_{\text{critical}}$ so the mind doesn’t shatter” — the ODE as the requirements spec for autopax (Gemini, AUDIT-WORKING-193847). The overwhelm/flooding analogy (large δ shutting down processing rather than scaling $\mathcal{T}$ up) is the same reach. Aspirational; routes to 04-e1i-core/ / autopax, not this segment.
- **Tempo arms-race / Red Queen reach.** Under symmetric adversarial coupling each agent’s tempo *becomes* the other’s volatility ($\rho_B = \rho_{B,\text{base}} + \gamma_A \mathcal{T}_A$), suggesting a mutual escalation that drives both toward their model-class fitness ceiling — a unified basis for arms races, Red Queen dynamics, economic competition (Gemini, AUDIT-WORKING-193847). Substantive home is `#result-adversarial-tempo-advantage` / the Part III adversarial segments.

Chapter 4

Persistence and Structural Limits

The cycle runs; the question is whether it can keep up. This chapter develops the central inequality of AAT — the persistence condition — together with what happens at and beyond its threshold: where it comes from, when the machinery can be expected to satisfy it, and what fails when it can't.

Chapter 3 ended with a preview. Under linear correction and bounded disturbance, the steady-state mismatch is $\rho/\mathcal{T}$ — disturbance over tempo. Mismatch stays bounded when the agent corrects faster than reality drifts. That was offered as a heuristic with a linear-ODE flavor. Chapter 4 makes it rigorous, generalizes it, and shows what it implies.

The first move is to replace “linear correction” with something more honest. Real correction mechanisms saturate at large mismatch, threshold at small mismatch, and break down entirely when the agent's model class is wrong for the situation. The sector condition — that the correction function points inward with at least baseline efficiency α within an operating region R — captures the qualitative essence of correction without committing to a specific functional form. This is one of the chapter's more satisfying pieces of theoretical hygiene: a condition flexible enough to admit saturating, sigmoid, threshold, and PID corrections under one Lyapunov argument, while still saying something concrete enough to derive a threshold from. Under that condition, the analysis gives

$$\alpha > \rho/R \tag{4.1}$$

three quantities, one threshold, mechanism underneath. The agent persists when its baseline correction efficiency exceeds the ratio of disturbance rate to operating reserve. Below the threshold, the containment guarantee is lost — escape from the operating region becomes possible, and for correctors whose sector bound is tight, certain under adversarial disturbance (A Lemma A.1N). The result is qualitative: persistence is not a degree, it's a regime — what disappears at the threshold is the certificate, not merely some margin. The inequality is genuinely domain-agnostic — its four variables map to biological adaptation (extinction is what happens when environmental change rate outpaces evolutionary correction), to organizational survival (the market shifts faster than the firm learns), to control instability, to cognitive overload. The framework identifies a structural pattern across these domains rather than a model of any one of them; whether the specific mechanism is the dominant cause in any given case is empirical, but the structure recurs because the math does.

The persistence condition has a thermodynamic shadow worth surfacing here, since the chapter derives it explicitly in I: under stochastic disturbance, maintaining the sector-persistence bound requires sustained Shannon information acquisition at rate $\dot{R} \geq n\alpha/2$ nats per unit time, with Kalman-Bucy saturating the bound. Survival, in this framework, is not a state you achieve once — it is a sustained burn rate. An agent cut off from informative observation begins to drift the moment the channel closes, and the rate of drift is the disturbance rate, independent of how good the agent was at adapting before. Two agents with identical persistence guarantees can face wildly different sustained demands. The threshold says whether persistence is *possible*; the information-rate corollary says what it *costs*.

But this isn't quite the operational form most domains care about. There is a second condition — task adequacy — that the steady-state mismatch be small enough for the agent's actions to remain useful. An agent can satisfy the structural condition (the machinery works) but fail task adequacy (the machinery doesn't work *well enough* for what the agent is trying to do). The two conditions have different remedies, and the split between them is the kind of move that's invisible until you see it stated and obvious in retrospect. A Kalman filter on a moving target satisfies structural persistence whenever it's stable, but task adequacy depends on what counts as "tracking" for the application. A software team can be structurally persistent but task-inadequate. The structural condition is what AAT derives; the task-adequacy threshold is a domain parameter — and conflating the two produces category errors in domain transfer.

The other move I want to flag is upstream. Where does α actually come from? Earlier in the framework, α shows up as a sector-condition parameter — a structural property of the correction function, taken as given. The gain-sector bridge (4.3) demotes it from postulate to derivation. For agents with directional fidelity in their update rule — including all optimal Bayesian updaters and all gradient agents on (locally) strongly convex losses — $\alpha = \eta^* \cdot c_{\min}$, where η^* is the gain principle from Chapter 3 and $c_{\min}$ is the worst-case directional fidelity of the update. The sector condition isn't a soft global postulate at all; for the cases AAT cares about most, it's a property of the gain. The five failure modes that break this derivation (directional infidelity, gain collapse, nonlinear saturation, unobservable directions, model misspecification) are named explicitly, and each connects back to a concrete mechanism rather than being left as residual error. This is the kind of move that makes the framework feel earned rather than assumed.

The chapter closes with two integrative results. *Structural adaptation necessity*: when model-class fitness $\mathcal{F}(\mathcal{M})$ is too low, no parametric update within the class can close the mismatch floor — the agent must change classes, not parameters. This is the consequence whose seed was planted in Chapter 2, now grown to a derived result. *Temporal nesting*: adaptive processes stratify by timescale, with each level operating on the quasi-steady-state output of the level below. Parametric update is fast; structural adaptation is slow; the convergence constraint between adjacent timescales means a slower process acting on transients of a faster one produces oscillation. This is standard singular-perturbation reasoning applied here to multi-level adaptation, but the consequence — that the timescale ratios are themselves constrained — is what licenses the abstraction throughout the rest of the framework.

A scope coda closes the chapter and Part I. AAT applies to agents on singular causal trajectories. The chronica $\mathcal{C}_t$ is not forkable — duplicating M_t and exposing the copies to different futures produces two agents on divergent trajectories, neither a sufficient statistic for the other's path. Identity in AAT lives in the trajectory, not in the state. This sits at the end of Chapter 4 because its dependencies (chronica from Chapter 1 + sufficiency from Chapter 2) span the section, but its substance is its own — and it's the bridge to the logogenic-agents and ELI work in Volumes 3 and 4, where the singular-trajectory commitment becomes load-bearing.

The flow of the chapter: deliberation cost (4.1) → gain-to-sector bridge (4.3) → sector-condition stability and the persistence inequality (4.4, 4.5) → structural adaptation when parametric update fails (4.8) → temporal nesting (4.9) → identity scope (4.10). The chapter is where Part I's machinery resolves into operational results.




Figure 3: Persistence as a bathtub. The belief–reality mismatch is the water level, drift is the inflow, correction is the drain, and model capacity is the rim. The same picture is drawn three times — literal water vocabulary, plain-English AAT terms, then formal symbols — in identical spatial positions, so the analogy is one-to-one. Overflow is persistence failure: inflow outpacing the drain at full.

- This is a chapter-introduction segment; it bridges Chapter 3’s preview of the central inequality to Chapter 4’s rigorous treatment and its consequences. It carries no formal claim of its own.
- Three moves the intro tries to surface as the chapter’s most interesting content: (1) the sector condition’s generalization of the linear case, preserving qualitative content without committing to functional form; (2) the structural-vs-task-adequacy split, which is non-obvious and load-bearing for domain transfer; (3) the gain-sector bridge, which derives α from the gain principle and demotes an opaque postulate to a property.
- The scope-coda paragraph addresses 4.10, which sits at the end of Part I Ch.4 because its dependencies span chapters. The intro names this rather than letting the segment appear unmotivated.
- The closing “where Part I’s machinery resolves into operational results” line is meant to register the chapter’s role as climax without overstating. The persistence condition is the single most-cited result in the framework downstream; Part I works because this chapter works.

Cross-audit “wandering thoughts” / §14-ideation lifted from the de-novo auditors’ working dirs (audit-routing-instructions.md §8), deduplicated across substrates and lightly attributed. Orthogonal pedagogical / framing material staged for a later Brief/Discussion-promotion pass — kept separate from certified theory-fix findings. Coverage on this intro: three substrates reached a digested reflection (Claude AUDIT-WORKING-384279; Codex/Claude AUDIT-WORKING-526815; Gemini AUDIT-WORKING-773921).

- The bathtub scaffold already in the segment was independently named the clearest possible physical intuition by the auditors, with the analog spelled out in convergent terms: water level = mismatch, inflow = drift ρ , drain = correction α , rim = capacity R , overflow = persistence failure (Gemini, AUDIT-WORKING-773921; Claude, AUDIT-WORKING-384279 — “the canonical respectful-pedagogy mental-model-first artifact”). Gemini extended the analog to *task adequacy* specifically: “the water is low enough that I can still see the bottom of the tub” — a candidate gloss the figure caption does not yet carry (Gemini, AUDIT-WORKING-773921).
- Survival-as-burn-rate framing for the thermodynamic shadow: “survival [is] not a state you achieve, but a sustained burn rate of information acquisition — if the channel closes, you die at rate ρ ” (Gemini, AUDIT-WORKING-773921). Tightens the existing thermodynamic-shadow paragraph toward a quotable opener.

- The field-contribution framing converged across substrates as the intro’s strongest one-liner: the persistence condition “unifies biological extinction, corporate bankruptcy, and control-system instability under a single geometric inequality” (Gemini, AUDIT-WORKING-773921). Already present in spirit in the domain-agnostic paragraph; the three-noun compression is sharper than the current prose.
- **Below-threshold phrasing (executed 2026-07-04).** Two substrates flagged “mismatch grows without effective bound” as over-stating the sector argument (Codex/Claude, AUDIT-WORKING-526815); the question they left open — does the downstream theorem actually prove escape? — was settled by the audit-731548 B-1 landing (#deriv-sector-condition Lemma A.1N, 2026-07-03: escape forced only in the radially tight case; general condition-failure means loss of the certificate). The body now carries the certificate-voice statement, consistent with #result-persistence-condition and #result-sector-condition-stability.
- **Branching chapter-roadmap diagram (not a theorem diagram).** The bathtub figure scaffolds the *math*; a separate, complementary figure would scaffold the *chapter*: linear ODE $\rightarrow$ sector generalization $\rightarrow$ (structural persistence + information cost) in one branch, task adequacy comparing bounded mismatch against a domain threshold in a parallel branch, with the path exiting to structural adaptation when bounded mismatch cannot be lowered inside the model class. The auditor’s explicit note: “this diagram should be a chapter map, not a theorem diagram” (Codex/Claude, AUDIT-WORKING-526815).

4.1 Derived: Deliberation Cost

The framework’s first formal trade-off between thinking and acting. Explicit deliberation can improve action quality by using the model for internal simulation before acting (improving the next update gain η^* — pausing praxis to improve upcoming epistrophe), but deliberation costs time, and during that time mismatch accumulates because the environment continues to evolve while the agent is not correcting (aporia accumulates during the pause). The derivation gives a clean threshold: deliberation of duration $\Delta\tau$ is net-beneficial when the gain improvement times post-deliberation mismatch magnitude exceeds the deliberation-window mismatch drift rate times the duration — stop deliberating when the marginal improvement rate drops below the mismatch drift rate normalized by the post-deliberation mismatch. Under diminishing returns (the first moments of simulation yield the largest gain improvement) combined with linear-in-duration drift cost, this yields a finite optimal deliberation duration; past that point, additional thinking is net-harmful.

The result is conditional on a *local deliberation-drift assumption* — that mismatch grows at an approximately constant rate ρ_{delib} during pause windows — which is a short-horizon assumption about inaction windows rather than a global dynamics model, weaker than the linear mismatch ODE (3.9), and can be estimated directly from pause windows in empirical traces. The Discussion below treats the consequences the framework reuses throughout the volume: high-drift environments penalize deliberation (the formal version of Boyd’s claim that over-deliberation is fatal in fast-tempo adversarial environments); implicit action as the high-tempo limit (formal grounding for why high-tempo environments favor action fluency); deliberation as investment when drift is low or post-deliberation mismatch is large; the self-consistency circularity (evaluating the threshold requires predicting post-deliberation mismatch using the same model deliberation is meant to improve); and the connection to temporal nesting as a singular-perturbation pattern.

Formal Expression.

Assumption (local deliberation drift):

During a deliberation pause of duration $\Delta\tau$, mismatch increases at an approximately constant local rate ρ_{delib} :

Assumption (deliberation-drift).

$$\Delta\|\delta\|_{\text{deliberation}} \approx \rho_{\text{delib}} \cdot \Delta\tau \quad (4.2)$$

This is a short-horizon assumption about inaction windows, not a full global dynamics model. It is weaker than the mismatch ODE and can be estimated directly from pause windows in empirical traces.

Proposition (deliberation threshold):

Deliberation of duration $\Delta\tau$ is net-beneficial when:

Derived (Conditional on deliberation-drift assumption).

$$\Delta\eta^*(\Delta\tau) \cdot \|\delta_{\text{post}}\| > \rho_{\text{delib}} \cdot \Delta\tau \quad (4.3)$$

where $\Delta\eta^*(\Delta\tau)$ is the improvement in post-deliberation update gain and $\|\delta_{\text{post}}\|$ is the mismatch magnitude the agent will face when it resumes acting.

Derivation.

1. Without deliberation, the agent acts immediately at current tempo $\mathcal{T}_0 = \nu \cdot \eta_0^*$.
2. With deliberation of duration $\Delta\tau$, the agent pauses, then acts with improved gain $\eta_0^* + \Delta\eta^*$. But during the pause, mismatch has grown by $\rho_{\text{delib}} \cdot \Delta\tau$.
3. The net mismatch reduction from acting after deliberation versus acting immediately: $\text{Net} = \Delta\eta^* \cdot \|\delta_{\text{post}}\| - \rho_{\text{delib}} \cdot \Delta\tau$.
4. Deliberation is justified iff $\text{Net} > 0$. $\square$

Optimal deliberation duration (under diminishing returns):

Derived (Conditional on diminishing-returns + deliberation-drift).

$$\Delta\tau^* = \arg \max_{\Delta\tau} [\Delta\eta^*(\Delta\tau) \cdot \|\delta_{\text{post}}\| - \rho_{\text{delib}} \cdot \Delta\tau] \quad (4.4)$$

where $\|\delta_{\text{post}}\|$ is treated as a parameter estimated by the agent before deliberation begins (not optimized over — the agent estimates the mismatch it will face, then decides how long to deliberate). Under this approximation, the first-order condition is: $\frac{\partial \Delta\eta^*}{\partial \Delta\tau} \cdot \|\delta_{\text{post}}\| = \rho_{\text{delib}}$. Stop deliberating when the marginal improvement rate drops below the mismatch drift rate (normalized by post-deliberation mismatch). When the dependence $\|\delta_{\text{post}}\| = \|\delta_0\| + \rho_{\text{delib}} \cdot \Delta\tau$ is included in the optimization, the exact FOC acquires a correction factor $(1 - \Delta\eta^*)$ on the cost side; this is negligible when $\Delta\eta^* \ll 1$ (the typical case — deliberation produces small gain improvements).

Epistemic Status.

Conditional on the local deliberation-drift assumption. The threshold condition is derived given the assumption; the assumption itself is a local approximation validated by consistency with the global mismatch dynamics (4.5). The result captures the *epistemic* benefit of deliberation (improving η^*); in practice, deliberation also provides a direct *action-value* benefit (choosing better actions that alter the environment trajectory), which operates through ρ reduction and immediate reward — a fuller formalization would incorporate the unified policy objective (3.7) at significantly more complexity.

Discussion.

High- ρ_{delib} environments penalize deliberation. When the environment changes rapidly during pause windows, the cost term grows quickly. Only very short deliberation with large $\Delta\eta^*$ can justify the pause. The model captures the same tradeoff Boyd emphasized: in fast-tempo adversarial environments, over-deliberation is fatal not because thinking is bad, but because the environment moves during the thinking. Whether the specific mechanism (mismatch drift during pause) is the dominant real-world effect is an empirical question.

Diminishing returns. In most models, $\Delta\eta^*(\Delta\tau)$ exhibits diminishing returns — the first moments of simulation yield the largest improvement. Combined with the linear cost $\rho_{\text{delib}} \cdot \Delta\tau$, this implies a finite optimal deliberation duration. Past that point, additional thinking is net-harmful.

Implicit action as the high-tempo limit. As $\rho_{\text{delib}} \rightarrow \infty$ or $\Delta\tau^* \rightarrow 0$: the optimal strategy converges to zero deliberation — pure implicit action (3.3). This provides a mathematical basis for why high-tempo environments favor action fluency: the cost of deliberation exceeds its benefit when $\Delta\eta^*$ is small (action-selection is already fluent) or ρ_{delib} is large.

Deliberation as an investment. When ρ_{delib} is low (stable environment) or $\|\delta_{\text{post}}\|$ is large (significant model-reality gap), deliberation pays off. The conditions favoring deliberation — stable environment, large mismatch — resemble the high-stakes, low-urgency scenarios where deliberative reasoning (System 2) is advantageous in dual-process theories. The structural parallel is suggestive; whether the cost-benefit mechanism is the same one governing System 1/System 2 selection is an open question.

The circularity of $\|\delta_{\text{post}}\|$. Evaluating the threshold requires the agent to *predict* post-deliberation mismatch using its current model — the same model deliberation is meant to improve. This circularity is typically benign: $\|\delta_{\text{post}}\|$ is bounded below by $\rho_{\text{delib}} \cdot \Delta\tau$ and above by current mismatch plus that accumulation. An agent that underestimates its mismatch will under-deliberate; one that overestimates will over-deliberate. The bias is self-correcting through the feedback loop. The threshold is best understood as a *design criterion*, not a real-time decision procedure.

Resource costs beyond time. Real agents also incur computational and energetic costs: internal simulation burns calories, compute cycles, or opportunity cost of not processing new observations. These are additive: $\Delta\eta^*(\Delta\tau) \cdot \|\delta_{\text{post}}\| > \rho_{\text{delib}} \cdot \Delta\tau + C(\Delta\tau)$. In high- ρ_{delib} environments the temporal cost dominates; in low- ρ_{delib} environments, resource costs may be the binding constraint.

Structural adaptation as an analogy. Structural adaptation (4.8) superficially resembles deliberation with a massive $\Delta\tau$: the agent's parametric loop is partially suspended while it searches for a new model class, incurring a large mismatch debt $\rho_{\text{delib}} \cdot \Delta\tau$. However, this is an informal analogy, not a consequence of the deliberation-cost formalism. Deliberation as formalized here improves η^* *within a fixed model class*; structural adaptation changes the model class itself, which is a mechanistically different operation (4.8). The cost-benefit structure may be similar in form, but the quantities involved ($\mathcal{F}(\mathcal{M})$ vs. η^* , model-class search vs. gain improvement) are distinct.

Connection to temporal nesting. Deliberation is a nested loop: internal simulation running at rate ν_{internal} within the external action loop at rate ν_{external} . The convergence constraint applies: the internal loop must approximately converge before the external loop acts on its output.

Connection to Part II. For actuated agents, the deliberation tradeoff extends to three modes: exploit (O_t via Σ_t), explore (improve M_t), and deliberate (revise Σ_t). The three-way allocation (10.2) extends this segment’s threshold by adding a strategic benefit term ΔV_2 ; the extended threshold is the one genuinely derived piece. The broader three-way framing is discussion-grade — simulation shows deliberation is rarely chosen by an oracle in simple settings, and a unified objective outperforms the two-stage decomposition.

The AI agent’s dilemma. An AI agent with 100% context turnover faces a severe version: it MUST deliberate (comprehend the codebase) before acting effectively, but during comprehension its context fills and the environment may change. The optimal comprehension depth depends on ρ_{delib} and the session’s action horizon. This is why reading CLAUDE.md and architecture docs first (high-CIY query actions) dominates reading random source files (low-CIY exploration).

Domain instantiations:

Table 4.1

<i>Domain</i>	<i>Deliberation</i>	$\Delta\eta^*$ <i>source</i>	<i>When ρ_{delib} is high</i>
Boyd’s OODA	Explicit “Decide” step	War-gaming, staff analysis	Collapses to IG&C (implicit)
RL / MCTS	Planning rollouts	Monte Carlo tree search	Fewer rollouts, shallower search
MPC	Online optimization	Trajectory optimization	Shorter horizons, faster solvers
Human cognition	System 2 deliberation	Mental simulation	Defaults to System 1 (intuition)
Organization	Strategic planning	Scenario analysis	“Move fast and break things”
Software developer	Reading code, analyzing alternatives	Architecture analysis	Ship now, refactor later
AI agent	Reading codebase, planning approach	Context-building	Limit comprehension, act sooner

Open questions:

1. *Computational cost of deliberation* is not just elapsed time but resource cost. A fuller model would include both temporal and computational budgets.
2. *Deliberation about deliberation*: deciding whether to deliberate itself takes time. This meta-deliberation is bounded by the same tradeoff at a higher level, suggesting a hierarchy of diminishing deliberation horizons.
3. *Deliberation that generates observations*: internal simulation can surface model inconsistencies (internal mismatch), functioning as “exploration without external action.” Can deliberation generate internal CIY?

Working Notes.

Incidental audit gold (2026-05-30 sweep). Cross-audit “wandering thoughts” / §14-ideation lifted from the de-novo auditors’ working dirs (audit-routing-instructions.md §8), deduplicated across substrates and lightly attributed. Orthogonal pedagogical / framing material staged for a later Brief/Discussion-promotion pass — kept separate from certified theory-fix findings. This was one of the most-reflected-on segments in Part I; coverage spans ten substrates (Gemini AUDIT-WORKING-193847/773921; Claude AUDIT-WORKING-266847/361742/384279/451729/471203/584721; Codex/Claude AUDIT-WORKING-526815/742613).

1. Candidate Brief prose / pre-prose.

- The deliberation threshold reads cleanly as a benefit-vs-cost ledger and the auditors converged on a plain-language gloss: left side is “the absolute benefit of having paused to think” (improvement in correction efficiency $\times$ size of error to correct), right side is “the cost of letting the world move while you were thinking” (rate of environment change $\times$ duration of pause); deliberate iff benefit $>$ cost (Gemini, AUDIT-WORKING-193847). A “pause ledger” framing — mismatch debt accumulating linearly during the pause, a mismatch-reduction credit booked after — is the same picture (Codex/Claude, AUDIT-WORKING-526815).

2. Candidate Discussion.

- **The action-bias / deliberation-bias regimes as a domain-spanning consequence.** The threshold “provides a formal justification for ‘action bias’ in startups/military and ‘deliberation bias’ in academia/planning, showing neither is universally correct — they optimize for different ρ_{delib} regimes” (Gemini, AUDIT-WORKING-193847); sharper still, “‘Move fast and break things’ is mathematically optimal when ρ_{delib} is high relative to $\Delta\eta^*$, and mathematically foolish when ρ_{delib} is low” (Gemini, AUDIT-WORKING-773921). The Discussion’s Boyd/System-2 material already gestures at this; the regime-as-optimization framing is a candidate sharpening.
- **Analysis-paralysis / perfectionism as a corollary of the FOC.** The first-order condition “proves that perfectionism is mathematically suboptimal in a changing universe: you must stop thinking and act the moment your marginal rate of insight drops below the speed of the world.” Analysis paralysis = overestimating $\Delta\eta^*$ or underestimating ρ_{delib} (Gemini, AUDIT-WORKING-193847). A vivid restatement of the existing diminishing-returns paragraph.
- **The benchmark/real-time asymmetry as a field observation.** “LLMs perform so well on static benchmarks (where time is frozen, $\rho = 0$, optimal strategy is infinite deliberation) but struggle in real-time continuous control — their architecture is tuned for $\rho = 0$ ” (Gemini, AUDIT-WORKING-193847). A striking diagnostic the segment could land for the logogenic-agents thread.

3. Follow-up items.

- **The “ $\Delta\eta^*$ vs action-value” scope gap — flagged by multiple substrates.** Two independent reads pressed that the prose opens by saying deliberation improves *action quality* via internal simulation, but the formal benefit term is improvement in *update gain* η^* — and many of the segment’s own examples (MCTS, MPC, war-gaming, reading code before editing) improve action selection / expected reward, not update gain. The Epistemic Status already names this, so neither counted it a finding, but both recommend the segment be *cited* as a threshold for “epistemic / gain-improving deliberation,” not all planning — one suggested retitling to “epistemic deliberation cost” and reserving the general term for a future action-value-bearing version (Codex/Claude, AUDIT-WORKING-526815; AUDIT-WORKING-742613). (*Routed as a scope-honesty follow-up; the strengthening direction — deriving the action-value benefit channel rather than deferring it — is noted in the report.*)
- **Derivation-prose tightness (the counterfactual baseline).** One auditor traced the “Net = $\Delta\eta^* \cdot \|\delta_{\text{post}}\| - \rho_{\text{delib}} \Delta\tau$ ” accounting and noted the implicit baseline (what the agent does *without* deliberating) is not made explicit: if “no deliberation” means continuous correction at tempo $\mathcal{T}_0$, the lost-correction term would read $\mathcal{T}_0 \|\delta\| \Delta\tau$ rather than $\rho_{\text{delib}} \Delta\tau$, unless ρ_{delib} is being read as an effective net-drift rate. Resolved as heuristic-but-consistent under the segment’s chosen rate-comparison framing (the conditional status and the “design criterion, not real-time procedure” caveat already cover it), but the prose “could be tightened to make the counterfactual and the exact accounting explicit” (Claude, AUDIT-WORKING-471203; the same initial concern was raised and then resolved-on-re-read by Claude, AUDIT-WORKING-584721). (*See report: candidate for tightening the derivation to an explicit case-A-vs-case-B accounting rather than a heuristic marginal comparison.*)
- **The $(1 - \Delta\eta^*)$ FOC correction factor.** Independently re-derived and confirmed by two substrates from $\|\delta_{\text{post}}\| = \|\delta_0\| + \rho_{\text{delib}} \Delta\tau$; both note the “ $(1 - \Delta\eta^*) \approx 1$ when $\Delta\eta^* \ll 1$ ” approximation is fine, and flag a watch-item: downstream uses of the threshold that omit the correction factor when $\Delta\eta^*$ is large would be a defect (Claude, AUDIT-WORKING-384279/584721).

- **“Open questions” inside a claims-verified segment.** One auditor noted the reader-facing “Open questions” block in a claims-verified segment “blurs the project convention slightly” (is this scope-honesty or stray Working Notes?) (AUDIT-WORKING-742613). Convention-texture signal, preserved rather than acted on.
- **Unified “search-during-pause” template (deliberation $\cup$ structural adaptation).** The segment correctly disclaims the structural-adaptation analogy as informal; one auditor wondered whether a genuine unified framing exists — both are “pause normal operation, search for a better state, resume,” differing only in *what space* is searched (deliberation searches gain space; structural adaptation searches model-class space) — and flagged it as a possible bigger-picture item, *not* a correction to the existing disclaimer (Claude, AUDIT-WORKING-584721).

4. Readers often ask / wonder.

- **Can deliberation raise $\mathcal{T}$ “for free”?** If an LLM runs internal chain-of-thought (a virtual $\mathcal{C}_t$) faster than physical action, does the internal generation let it raise effective tempo without paying the full ρ_{delib} penalty? (Gemini, AUDIT-WORKING-193847). Connects to open-question 3.
- **Where does the meta-deliberation regress bottom out?** Open-question 2 (“deliberation about deliberation”) is the framework’s first brush with metacognition; readers will ask whether the hierarchy of diminishing deliberation horizons terminates, and the assertion that meta-costs dominate at each level is currently un-derived (Claude, AUDIT-WORKING-266847; Gemini, AUDIT-WORKING-773921 — “AAT seems to dodge the infinite regress by treating the threshold as a design criterion, not a real-time computation step”).

5. Candidate figures.

- **Two-channel benefit diagram.** A diagram showing the formal benefit channel ($\Delta\eta^*$, gain improvement) in solid lines and the acknowledged-but-not-derived action-value channel in dashed lines — making the scope boundary visual rather than buried in Epistemic Status (Codex/Claude, AUDIT-WORKING-526815). Pairs directly with the $\Delta\eta^*$ -vs-action-value follow-up above.

6. Belongs elsewhere.

- **ELI session-start instantiation (Section IV / 04-eli-core/).** The “AI agent’s dilemma” generalizes to the developmental case: an ELI’s session begins with high U_M (no context), and the criterion says comprehension should continue until $\Delta\eta^* \cdot \|\delta_{\text{post}}\| \approx \rho_{\text{delib}}\Delta\tau$ — until the marginal value of more reading equals the cost of context-filling. For a hostile / high- ρ environment the infrastructure must enforce a bias toward action so the agent does not “overthink to death” (Claude, AUDIT-WORKING-471203; Gemini, AUDIT-WORKING-193847). Aspirational reach pointing at the ELI work, not at this segment.
- **CIY over-exploration failure (adjacent to #def-causal-information-yield).** A consequence of the CIY-vs-EIG distinction surfaced while reading this segment: an agent exploring “because actions are distinguishable” keeps acting even after the causal graph is fully characterized — high CIY, zero expected information gain; the $\lambda(M_t)$ weighting is a heuristic patch, a proper EIG formulation the clean fix. The code-work analog: an agent that keeps running tests after the bug is fixed (Claude, AUDIT-WORKING-451729). Belongs with the CIY segment, noted here only because it co-occurred in the deliberation sweep.
- **Naming seed.** “Deliberation Threshold” / “Think-vs-Act Tradeoff” floated as a segment-title alternative that surfaces the operational content faster than “deliberation cost”; the AAT-distinctive content is the *threshold*, not the term (Claude, AUDIT-WORKING-471203). Naming-cycle target, not a segment edit.

4.2 Formulation: Sector Condition (A2')

The sector condition (A2') is the structural shape AAT chooses for the correction-function geometry: the correction points roughly inward, with magnitude bounded below relative to the mismatch, on a local region of validity. Together with the companion structural properties (A1) zero-correction-at-zero-mismatch and (A3) tempo-monotonicity, A2'

is the formal expression of “the agent’s adaptation tracks reality with at least baseline efficiency” — the form on which the persistence and adaptive-reserve results of A rest, and the form which downstream consumers (`#deriv-discrete-sector-condition`, `#deriv-variational-sector-condition`, `#deriv-adaptive-gain-dynamics`, `#der-gain-sector-bridge`) extend, weaken, or ground.

Formal Expression.

Setup objects (carried into the form). Let $\delta(t) \in \mathbb{R}^n$ be the mismatch vector (3.4 — the difference between the model’s predictions and reality across n observable dimensions). Let $F(\mathcal{T}, \delta) : \mathbb{R}_+ \times \mathbb{R}^n \rightarrow \mathbb{R}^n$ be the **correction function** — how the agent’s adaptive process reduces mismatch — mapping into the same space as δ so that the inner product $\delta^T F$ is well-defined. This subsumes the update gain η^* (3.6), event rate ν , and the structure of the update rule. The adaptive tempo $\mathcal{T}$ (3.8) is the rate parameter.

(A1) Zero Correction at Zero Mismatch.

Assumption A1.

$$F(\mathcal{T}, 0) = 0 \quad (4.5)$$

No correction is applied when the model perfectly matches reality. Uncontroversial by construction.

(A2’) Local Sector Condition. There exists a region $\mathcal{B}_R = \{\delta : \|\delta\| \leq R\}$ and $\alpha > 0$ such that (following the sector-condition framework of Lur’e¹):

Formulation A2’ (sector-condition) — derived in sub-scope α , assumed in sub-scope β (see Grounding below).

$$\delta^T F(\mathcal{T}, \delta) \geq \alpha \|\delta\|^2 \quad \forall \delta \in \mathcal{B}_R \quad (4.6)$$

The correction function always points “inward” (reducing mismatch), and its magnitude is bounded below relative to $\|\delta\|^2$. The linear case has $\alpha = \mathcal{T}$. A saturating correction has α decreasing for large $\|\delta\|$. A threshold correction has $\alpha = 0$ for small $\|\delta\|$.

The local form allows the correction to break down outside $\mathcal{B}_R$ — the structural adaptation regime of 4.8.

(A3) Tempo Monotonicity.

For fixed δ , $\delta^T F(\mathcal{T}, \delta)$ is monotone increasing in $\mathcal{T}$. Higher tempo means faster correction.

Assumption A3.

Parameter interpretation. The sector parameter α is determined by the adaptive tempo $\mathcal{T}$ and the structure of the correction function. In the linear case, $\alpha = \mathcal{T} = \sum_k \nu^{(k)} \cdot \eta^{(k)*}$. In nonlinear cases, α represents the *worst-case* correction efficiency within the valid region — the minimum ratio of correction power to mismatch magnitude. The radius R represents the model class capacity: how large a mismatch can grow before the correction mechanism fails (i.e., before the sector condition ceases to hold), at which point structural adaptation (4.8) becomes necessary.

¹Lur’e, A. I. (1957). *Some Nonlinear Problems in the Theory of Automatic Control*.

Epistemic Status.

A2' is a *formulation* — a chosen structural shape for the correction-function geometry that captures “the correction points roughly inward with bounded efficiency.” Several structurally distinct shapes would fit the same role (a sign-only condition; a two-point monotonicity; a contraction-rate bound; a non-Euclidean weighted inner-product variant — see *Why Euclidean A2' specifically* in Discussion below). A2' is the canonical choice *matched to the quadratic Lyapunov candidate* $V = \frac{1}{2}\|\delta\|^2$ used in A, and matched to the gain-based update form on which 4.3 derives it for sub-scope α .

The status is *conditional* because A2' is **derived** for one explicitly named sub-scope of AAT-in-scope agents and **assumed as a per-system empirical claim** for the complementary sub-scope. The Lyapunov derivations downstream of A2' (A Props A.1, A.1S, A.2; Corollary A.1S.1) apply uniformly across both sub-scopes — they operate downstream of A2' regardless of how it is established. The sub-scope partition is *scope narrowing*, not *scope retreat*: it makes explicit which agent classes get A2' as a derived consequence of their update geometry, and which classes carry it as a structurally well-scoped posit.

Sub-scope α (A2' derived). For a characterized class of AAT-in-scope agents, A2' is a *derived* consequence of the update rule, not a primitive assumption. 4.3 (Prop B.3) shows that the gain-based update $M_t = M_{t-1} + \eta^* g(\delta_t)$ (3.6) induces a correction function satisfying A2' whenever the update rule has **directional fidelity (B1)** — $\delta^T H g(\delta) \geq c_{\min} \|\delta\|^2$ on $\mathcal{B}_R$. Sub-scope α — the agent classes where B1 holds structurally — includes:

- *Optimal Bayesian updates* (Kalman, conjugate families): B1 holds by Bayes-risk minimization. $\alpha = \eta^* \cdot c_{\min}$, reducing to $\alpha = \eta^*$ in the scalar case.
- *Exponential families in natural parameters*, on a bounded interior scope $\Theta_0 \subset \text{int}(\Theta)$: the Hessian is the Fisher information matrix, PD on the interior. $\alpha = \eta \cdot \mu_0$ where $\mu_0 = \inf_{\theta \in \Theta_0} \lambda_{\min}(\mathbf{I}(\theta)) > 0$. Pointwise PD does not imply a uniform global lower bound: $\inf_{\theta \in \Theta} \lambda_{\min}(\mathbf{I}(\theta))$ can be zero (Poisson: $\mathbf{I}(\theta) = e^\theta$, $\inf_{\theta \in \mathbb{R}} e^\theta = 0$), so Θ_0 supplies the local-region scope R that A2' requires. Families with a uniform Fisher floor on Θ (Gaussian-mean, Beta-Bernoulli) extend to global α .
- *Gradient descent on locally strongly convex losses* (Prop B.4): B1 is *equivalent* to strong convexity via the gradient-monotonicity characterization (Nesterov 2004², Thm 2.1.10). $\alpha = \eta \cdot \mu$ where μ is the strong convexity modulus.
- *L2-regularized convex losses*: regularization provides a floor $\mu \geq \lambda$, so $\alpha \geq \eta \lambda$ globally.
- *Linear corrections with positive-definite gain-observation product*: $\alpha = \lambda_{\min}^+(KH)$ (matrix Kalman, restricted to the observable subspace).

Within sub-scope α , A2' is written down by inspection of the update rule; no independent posit is required.

Sub-scope β (A2' assumed as empirical claim). For the remaining AAT-in-scope agents, A2' stands as a well-scoped empirical claim about the agent's correction geometry. Sub-scope β includes:

- *PID controllers with fixed gains* — no gradient / optimality structure; B1 is a tuning question, not a structural consequence.
- *Rule-based systems* — no continuous update rule; A2' is domain-specific (see also the structural Lipschitz-floor scope-exit in Discussion below).
- *Human judgment / organizational learning* — structural analogy in 3.6; no formal B1 guarantee.
- *Severely misspecified agents* (FM-5 in 4.3) — proper-gradient rules can aim at the wrong target.

²Nesterov, Y. (2004). *Introductory Lectures on Convex Optimization*. Springer. Theorem 2.1.10.

- *Variational / approximate posteriors* — B1 not guaranteed by optimality because approximation error can rotate the correction. Under a controlled KL bound, an intermediate sub-scope α' recovers a degraded form of A2' (see AJ).
- *Non-convex gradient agents beyond the basin* — A2' fails at basin boundary; this IS the 4.8 trigger.
- *Stochastic gradients (per-step)* — A2' holds in expectation; per-step noise enters as effective disturbance under A Prop A.1S.

Within sub-scope β , A2' must be verified per-system — the claim is stronger than what AAT's postulates + gain structure alone can force (an agent with $g(\delta) = R_{90^\circ}\delta$ satisfies every AAT postulate but violates B1 trivially; see 4.3 FM-1).

Discussion.

Operator-family classification (external mathematical lineage). The sub-scope α/β partition can be restated precisely in the operator-theoretic language of Rockafellar 1970 (*Convex Analysis*, §24) and Bauschke & Combettes 2017 (*Convex Analysis and Monotone Operator Theory in Hilbert Spaces*, 2nd ed., §§22–28). Casting the discrete update map as $T_d(\delta) = \delta - \eta^* F_d(\delta)$ with $T_d(0) = 0$, A2' at $\delta^* = 0$ is exactly a **one-point strong monotonicity condition** on $A = I - T_d$, and DA2'-inc (the incremental strengthening required by #form-composition-closure's bridge lemma, stated in ○) is **full two-point strong monotonicity** in the Rockafellar sense. Sub-scope α maps onto established operator families:

- **Optimal Bayesian updates** (Kalman, conjugate, exponential family) $\equiv$ proximal / firmly nonexpansive operators (Bauschke-Combettes Prop 23.7). Firmly nonexpansive is equivalent to $\frac{1}{2}$ -averaged; operator-sector condition holds with the gain as the sector constant.
- **Gradient operators on strongly convex functions** $\equiv$ cocoercive operators via the Baillon-Haddad theorem; strong monotonicity with modulus $\eta\mu$.
- **Exponential families in natural parameters** $\equiv$ natural-gradient operators in the Fisher-weighted Hilbert space (Amari 1998; cf. 4.3 “Fisher-metric cases under parameterization-invariance” for the AAT-internal forcing via (PI)/Čencov named in 5.5).
- **L2-regularized convex gradients** $\equiv$ shifted monotone-gradient operators with regularization-provided floor.
- **Linear corrections with PD gain–observation product** $\equiv$ linear contraction in the $(P^-)^{-1}$ -weighted inner product.

Sub-scope β maps onto “operator families that are not cocoercive in any natural inner product”: PID (tuning-dependent; not cocoercive generically); rule-based / symbolic (no inner-product structure); variational / amortized (approximation-gap rotation); non-convex-gradient beyond basin (non-monotone on the full domain); per-step SGD (noise-dominated); human judgment (no formal rule). The α/β epistemic labeling in AAT tracks whether the operator family admits operator-sector structurally (derived from the class definition) or only under empirical per-instance verification — a scope-honesty move, not a mathematical reclassification.

This recognition positions AAT's sector-condition framework as a specialization of monotone-operator theory (Minty 1962; Browder 1968; Rockafellar 1970; Bauschke-Combettes 2017) to one-point-anchored strong monotonicity under a specific inner product structure. AAT's distinctive content — one-point anchoring at the equilibrium (strictly weaker than full two-point strong monotonicity, matched to fixed-point-at-target semantics); Model D / Model S disturbance decomposition; composition with the identifiability-floor meta-pattern; composition-consistency postulate; the α/β epistemic labeling itself — sits as specialization + repurposing rather than strict generalization. This acknowledgment is load-bearing for scope honesty: the mathematical machinery is established; AAT's value lies in its AAT-internal architecture (organization-of-scope under a singular-trajectory agent, signed-coupling structure, forced coordinates via uniqueness theorems, three meta-patterns) rather than in novel monotone-operator mathematics. The unification has

honest limits: the coarse-graining projection Λ (#form-composition-closure) does not fit the operator-sector primitive; three of five metric- α_2 cases in #result-contraction-template remain theorem-imported; the identifiability-floor axis is orthogonal.

Sub-scope β “rule-based / discontinuous” — structural Lipschitz floor. [Derived, status: exact scope-exit statement.] The rule-based / state-machine / threshold-triggered / discontinuous entry in sub-scope β is not merely “empirical verification required”; it is a **structural scope-exit for contraction-based bridge-lemma analysis**. For correction functions F that are not locally Lipschitz, no scalar sector bound $\delta^T F(\delta) \geq \alpha \|\delta\|^2$ implies the full-update-map contraction required by #form-composition-closure’s bridge lemma — the update map $f_c(X, o) = X - \eta^* F(X)$ can exhibit $\Omega(1)$ jumps between arbitrarily close X, X' at rule-firing boundaries, so no $\lambda < 1$ Lipschitz constant exists. Concrete counterexample: $F(\delta) = \alpha\delta$ for $|\delta| < 1$ and $F(\delta) = \alpha\delta + \text{sign}(\delta)$ for $|\delta| \geq 1$ — the sector bound holds with α , but f_c jumps by η^* at $|X| = 1$, violating contraction. The appropriate external apparatus for this class is the **hybrid-dissipative framework** (van der Schaft & Schumacher 2000, *An Introduction to Hybrid Dynamical Systems*, Springer; Di Bernardo, Liuzza & Russo 2014, *SIAM J. Control Optim.* 52) — distinct machinery (dwell-time, Filippov solutions, impulsive-dissipative inequalities) operating on a different regularity class. AAT’s Lyapunov machinery in A and the contraction machinery in #result-contraction-template both require C^1 or better regularity; the scope-exit to hybrid-dissipative analysis is honest and structural, not a deficiency to be repaired within AAT’s current apparatus. Under #disc-identifiability-floor’s Instance 2 pattern (Cramér-Rao floor under unobservable common cause), rule-based agents whose rule firing depends on regime structure — e.g., threshold-triggered state-machines with regime-dependent thresholds — additionally suffer regime-C identifiability collapse when the regime variable is unobservable; the non-contraction and the non-identifiability compose in that case.

Why Euclidean A2' specifically. The A2' form $\delta^T F \geq \alpha \|\delta\|^2$ is the sector condition *matched to the quadratic Lyapunov candidate* $V = \frac{1}{2} \|\delta\|^2$ used by A. A converse-Lyapunov argument (Khalil 2002³, Thm 4.17) gives: if persistence holds under the dynamics $\dot{\delta} = -F(\delta)$ on $\mathcal{B}_R$, then there exists a quadratic-equivalent Lyapunov function $V_*(\delta)$ with $c_1 \|\delta\|^2 \leq V_* \leq c_2 \|\delta\|^2$ — but V_* may not be the Euclidean norm itself. Under a weighted Lyapunov candidate $V(\delta) = \frac{1}{2} \delta^T P \delta$, the natural sector condition is $\delta^T P F(\delta) \geq \alpha \delta^T P \delta$; the matrix-Kalman case of 4.3 is exactly this in the $(P^-)^{-1}$ -weighted inner product, with a norm-equivalence transfer to Euclidean A2' degraded by the condition number $\kappa(P^-)$. Euclidean A2' is therefore not the unique sector form — it is the canonical one matched to the canonical V . An agent that persists under a non-Euclidean metric satisfies a weighted-sector A2' that transfers to Euclidean A2' only up to norm equivalence. 5.5 classifies the Lyapunov case as an *adjacent family member* to AAT’s three primary additive-coordinate-forcing instances (chain / divergence / update, where a logarithmic coordinate is uniquely forced via Cauchy’s functional equation under an AAT-internal additivity axiom): here the quadratic coordinate is matched to the sector form rather than forced by one.

How downstream segments use this form.

- A consumes A2' as the Lyapunov-derivation input: Props A.1 / A.1S / A.2 prove ultimate boundedness, mean-square persistence, and adaptive reserve under (A1), (A2'), (A3); Corollary A.1S.1 establishes the disturbance-model containment dichotomy $P(\tau_R < \infty) \in \{0, 1\}$.
- B demonstrates the load-bearing Model-S half of Cor A.1S.1 — that no horizon-independent non-exit bound exists under additive stochastic forcing — using A2' on $\mathcal{B}_R$ as the Itô-Lyapunov input.
- 4.3 derives A2' for sub-scope α under directional fidelity (B1), making the sub-scope- α partition above structural rather than postulated.
- O states the discrete-time analog (DA2'), strengthening A2' with an additional Lipschitz bound on $\|F_d\|$ to control the quadratic term that arises in discrete contraction.

³Khalil, H. K. (2002). *Nonlinear Systems* (3rd ed.). Prentice Hall.

- **AJ** states ε -fidelity $A2'$ under controlled-KL variational approximation, with $O(\sqrt{\varepsilon})$ sector-constant degradation — recovering an intermediate sub-scope α' within the partition above.
- **V** refines sub-scope α into fixed-gain (α_1) and adaptive-gain (α_2 under meta-gain conditions MG-1–MG-4) layers.
- **4.4** and **G** state the persistence results in result-segment voice; **13.1** and the composite-agent segments in Part III consume $DA2'$ -inc as the bridge-lemma precondition.

Working Notes.

- Landing-context provenance: this segment was carved out of **A** in the 451729 D.1 Gate verification cycle (2026-05-20) to resolve a depends-graph topology issue — **4.3** and **B** both already depended on **A**, so making the Lyapunov-derivation segment depend on its own $A2'$ -form-consumers would have created cycles. The sub-scope α/β partition reflects the strengthening trail recorded in `spikes/.integrated/spike-a2-prime-strengthening.md` — the analysis that ruled out a universal $A2'$ derivation and identified the five operator families where the bridge is structural.
- The companion-property partition ($A1, A2', A3$ together) is held here as the canonical setup; alternate orderings or sub-scope refinements (e.g., the $\alpha_1/\alpha_2/\alpha_3/\alpha'$ ladder accumulated across `#deriv-variational-sector-condition`, `#deriv-adaptive-gain-dynamics`, `#deriv-fisher-whitened-update-rule`) live in their respective segments and inherit this form by reference rather than re-stating it.

Incidental audit gold (2026-05-30 sweep). Cross-audit “wandering thoughts” / §14-ideation lifted from the de-novo auditors’ working dirs (`audit-routing-instructions.md` §8). Orthogonal framing material staged for a later Brief/Discussion-promotion pass — kept separate from certified theory-fix findings. Thin coverage on this segment (one digested reflection reached it: Claude, `AUDIT-WORKING-384279`); the gold here is methodological-positioning rather than pedagogical, but worth preserving.

2. Candidate Discussion.

- The auditor read this segment as the framework’s *methodological-honesty exemplar* — the operator-family classification (mapping the sub-scope α derivation forms to Rockafellar / Bauschke–Combettes 2017 / Baillon–Haddad) “names AAT’s reach into monotone-operator theory with humility ... AAT’s distinctive content sits as specialization + repurposing rather than strict generalization.” Candidate framing-prose for *why* the sub-scope partition is structured the way it is: it is the template for how scope-exit and prior-art positioning should be done (Claude, `AUDIT-WORKING-384279`).
- The **Lipschitz-floor scope-exit** (sub-scope β rule-based / discontinuous correction) was singled out as exemplary: it is not merely “verify per-system” but a *structural* scope-exit for contraction-based bridge-lemma analysis, with a concrete counterexample and the right external apparatus (van der Schaft & Schumacher 2000, hybrid-dissipative). Flagged as a model other segments’ scope-exits should imitate (Claude, `AUDIT-WORKING-384279`).

3. Follow-up items.

- **Respect the $\kappa(P^-)$ caveat downstream.** The condition-number caveat on the matrix-Kalman $\rightarrow$ Euclidean- $A2'$ transfer (“degraded by the condition number”) was confirmed correct; the watch-item is that downstream segments invoking $A2'$ via this transfer should carry the caveat rather than assuming the clean Euclidean form (Claude, `AUDIT-WORKING-384279`).
-

4.3 Derived: Gain–Sector Bridge

The gain-based update principle (3.6) produces correction dynamics satisfying the sector condition (GA-3) whenever the update rule has *directional fidelity* — the correction points at least roughly toward reality. For gradient-based agents, local strong convexity of the loss is sufficient for the one-point sector condition ($A2'$ as stated in 4.2) and bidirectionally equivalent to the two-point / incremental sector condition ($DA2'$ -inc). The sector parameter α is not a free parameter but is determined by the gain and the correction geometry.

Formal Expression.**The Bridge Theorem.**

Given the gain-based update $M_t = M_{t-1} + \eta^* \cdot g(\delta_t)$ (3.6), the induced correction function is:

Derived (gain-sector-bridge, from update-gain + directional fidelity).

$$F(\delta) = \eta^* \cdot H g(\delta) \quad (4.7)$$

where H maps state-space corrections to observation-space mismatch reduction. The sector condition (GA-3) holds with parameter $\alpha > 0$ whenever:

(B1) Directional fidelity. The mismatch transform g preserves the mismatch-reducing direction:

$$\delta^T H g(\delta) \geq c \|\delta\|^2 \quad \text{for } \|\delta\| \leq R \quad (4.8)$$

for some $c > 0$. The bridge then determines the sector parameter in two forms that differ by the event-rate normalization, because the gain-based update is applied *per correction event* while the sector condition feeds a *continuous-time* persistence inequality.

Per-event sector efficiency — the dimensionless fraction of standing mismatch the correction removes per event:

$$\alpha_{\text{event}} = \eta^* \cdot c_{\min}, \quad c_{\min} = \inf_{\|\delta\| \leq R} \frac{\delta^T H g(\delta)}{\|\delta\|^2} \quad (4.9)$$

with η^* dimensionless and $c_{\min}$ a dimensionless geometric (Rayleigh-quotient) ratio, so $\alpha_{\text{event}} \in (0, 1]$ is itself dimensionless.

Per-time sector rate — the correction *rate* the continuous-time sector condition (GA-3) and the persistence inequality $\alpha > \rho/R$ actually compare against, obtained by aggregating per-event efficiency over the channel's event rate ν (3.8):

$$\alpha = \alpha_{\text{time}} = \nu \cdot \eta^* \cdot c_{\min} \quad (4.10)$$

with units t^{-1} — the same units as adaptive tempo $\mathcal{T}$ (NOTATION.md *Units*). The event rate ν is exactly the factor that promotes the per-event geometric quantity to the per-time rate that the Lyapunov machinery in A consumes; the bare α in GA-3 and downstream is α_{time} .

Gradient Equivalence.

The sector parameters in this section and the Verified-Instances table below are stated in the *per-event* form α_{event} (the dimensionless correction efficiency); each promotes to the per-time rate the persistence inequality consumes by the same factor ν , $\alpha_{\text{time}} = \nu \cdot \alpha_{\text{event}}$, so the proportionalities are identical under either reading.

Derived (sector-convexity equivalence, two-point form).

For any agent updating via gradient descent on a loss L with learning rate η :

$$\alpha = \eta \cdot \mu \quad \text{where } \mu = \inf_{\|\delta\| \leq R} \lambda_{\min}(\nabla^2 L(M^* + \delta)) \quad (4.11)$$

is the strong convexity modulus. The basin radius R is the largest ball around the optimum where $\nabla^2 L$ remains positive definite. The equivalence has two forms — bidirectional under the stronger two-point sector condition, one-directional under the one-point form actually used by `#deriv-sector-condition`:

- **Two-point / incremental sector $\Leftrightarrow$ strong convexity (full equivalence).** Under the incremental sector bound $(F(\delta_1) - F(\delta_2))^T(\delta_1 - \delta_2) \geq \alpha \|\delta_1 - \delta_2\|^2$ on $\mathcal{B}_R(M^*)$ — DA2'-inc in O, the bridge-lemma precondition in 13.1 — the iff holds via Nesterov 2004 Theorem 2.1.10: Two-point sector with $(\alpha, R) \iff L$ is locally (α/η) -strongly convex on $\mathcal{B}_R(M^*)$.
- **One-point sector $\Leftarrow$ strong convexity (one direction only).** AAT's GA-3 / A2' as stated in 4.2 is the one-point form $\delta^T F(\delta) \geq \alpha \|\delta\|^2$ at $\delta^* = 0$. Strong convexity implies the one-point sector ($\alpha = \eta\mu$); the converse fails. Counterexample: $L'(x) = x(1 + \frac{1}{2} \sin(10x))$ satisfies $x \cdot L'(x) \geq \frac{1}{2}x^2$ globally yet has $L''(\pi/10) < 0$, so it is not convex on any neighborhood of $x^* = 0$. The one-point sector at the equilibrium is genuinely weaker than full local strong convexity (cf. G's one-point/two-point distinction). Full proofs and the counterexample analysis in K Prop B.4.

Verified Instances.

Table 4.2

Update class	Bridge status	Sector parameter α	Valid region
Scalar Kalman	Derived	$K = P^-(P^- + R_{\text{obs}}) = \eta^*$	Global
Matrix Kalman	Derived	$\lambda_{\min}^+(KH)$ in $(P^-)^{-1}$ -norm	Observable subspace
Beta-Bernoulli	Derived	$1/(n+1) = \eta_{\text{edge}}$	Global
Exponential family (natural params), bounded scope $\Theta_0 \subset \text{int}(\Theta)$	Derived	$\eta \cdot \mu_0$ where $\mu_0 = \inf_{\theta \in \Theta_0} \lambda_{\min}(\mathbf{I}(\theta)) > 0$	Θ_0 (compact / interior-bounded) — global only when the family has a uniform Fisher lower bound
Gradient on strongly convex loss	Derived	$\eta \cdot \mu$	Global ($R = \infty$)
Gradient on locally convex loss	Derived	$\eta \cdot \mu_{\text{local}}$	Basin of attraction
Gradient on non-convex loss	Fails at basin boundary	N/A beyond R	Finite R
SPR-tuned PID on positive-real plant with anti-windup	Derived	$\alpha_{\text{PID}} = \omega_c \sin(\varphi_m)/\kappa(P)$ (phase margin as sector constant; crossover frequency as tempo; KYP-certificate condition number as degradation)	Classical linear regime + Lur'e sector-bounded nonlinearity within specified plant-Lipschitz threshold

Failure Modes. The bridge fails precisely in five cases:

1. **Directional infidelity.** The mismatch transform g rotates the correction away from the mismatch ($\delta^T Hg(\delta) \leq 0$). Occurs with pathological parameterizations or severe model-observation misalignment. For optimal Bayesian updates, B1 holds by construction.
2. **Gain collapse.** $\eta^* \rightarrow 0$ while $\rho > 0$, so $\alpha \rightarrow 0$ and the persistence condition eventually fails. Not a failure of the bridge but of the persistence condition — see the gain-collapse analysis in 3.6.
3. **Nonlinear saturation.** The correction function g saturates at large $\|\delta\|$, so the sector ratio $\delta^T g(\delta)/\|\delta\|^2$ decays. The sector condition holds locally with α depending on R . This is exactly what A2' (the local sector condition) is designed for.
4. **Unobservable directions.** When $\ker(H) \neq \{0\}$, the correction has no effect on mismatch in unobservable directions. The sector condition holds only in the observable subspace. See 9.3.

5. **Model misspecification.** The model class does not contain the truth, so the gradient direction is wrong. B1 fails because the correction aims at the wrong target. This is the 4.8 trigger.

Epistemic Status.

Conditional derivation. The bridge theorem is exact for all cases where B1 (directional fidelity) holds. The resulting sub-scope structure of A2' in 4.2 is:

Sub-scope α (B1 structural, A2's derived):

- **Optimal Bayesian updates** (Kalman, conjugate, exponential family): B1 holds by construction — the posterior update minimizes expected loss, ensuring the correction aligns with the mismatch. The sector parameter equals the gain: $\alpha = \eta^*$ (scalar) or $\alpha = \lambda_{\min}^+(KH)$ (matrix Kalman, observable subspace).
- **Gradient descent on (locally) strongly convex losses:** local strong convexity is *sufficient* for B1 (the one-point form of A2' at M^*) and *bidirectionally equivalent* to the two-point / incremental sector condition DA2'-inc (Prop B.4 (B.4-i) and (B.4-ii) respectively) — a well-characterized property with an extensive optimization theory literature. The sector parameter factors as $\alpha = \eta \cdot \mu$ (learning rate $\times$ curvature). The reverse implication B1 $\Rightarrow$ strong convexity does *not* hold without the two-point upgrade: the one-point sector at M^* is strictly weaker than full local strong convexity (counterexample in Prop B.4).
- **L2-regularized convex losses:** the regularization parameter λ provides a global floor $\mu \geq \lambda$, so $\alpha \geq \eta\lambda$ globally.
- **Exponential families in natural parameters, on a bounded interior scope:** Fisher information matrix is PD on the interior, and uniformly bounded below on any compact $\Theta_0 \subset \text{int}(\Theta)$; $\alpha = \eta \cdot \mu_0$ with $\mu_0 = \inf_{\theta \in \Theta_0} \lambda_{\min}(\mathbf{I}(\theta))$. The bound is global only when the family has a uniform Fisher lower bound on Θ — true for Gaussian-mean and Beta-Bernoulli, false for Poisson natural parameter ($\mathbf{I}(\theta) = e^\theta$, infimum zero).
- **Linear corrections with PD gain–observation product:** $\alpha = \lambda_{\min}^+(KH)$.

Within sub-scope α , A2' is written down by inspection of the update rule — no independent posit is required. This is what #form-sector-condition “Sub-scope α (A2' derived)” names.

Sub-scope β (B1 not structural, A2's assumed per-agent):

- **Non-gradient agents** (PID controllers, rule-based systems, human judgment): B1 remains an empirical claim. Well-tuned PID has B1 empirically; badly-tuned PID may violate it.
- **Severely misspecified agents** (FM-5 below): proper-gradient updates can aim at the wrong target.
- **Variational / approximate posteriors:** B1 not guaranteed by optimality — approximation-direction error can rotate the correction.
- **Non-convex gradient agents beyond the basin** (FM-3 + basin boundary): A2' fails where the loss curvature goes non-positive; the structural-adaptation-necessity trigger.
- **Stochastic gradients, per-step:** A2' holds in expectation; per-step noise enters as effective disturbance under Prop A.1S.

The bridge covers sub-scope α rigorously and characterizes the boundary to sub-scope β via the five failure modes. It does *not* eliminate GA-3 as an assumption for all AAT-in-scope agents — some agent classes genuinely require A2' as a primitive posit, and the honest architectural statement is scope narrowing rather than universal derivation.

The gradient equivalence is validated by simulation across quadratic, logistic, exponential-family, and non-convex losses. The Kalman case is verified analytically. Full derivations and simulation results in [K](#).

Max attainable: conditional — the condition (B1 or strong convexity) is inherent, not removable. Pathological update rules exist that violate B1 (FM-1 provides a counterexample that satisfies every AAT postulate and the gain-based update form but has $\delta^T g(\delta) = 0$ identically). Scope + gain structure alone does not force B1; some optimality / coherence / rationality constraint (Bayesian coherence, gradient-of-a-convex-loss, etc.) is required.

Weighted-norm subtlety. In the matrix Kalman case, the sector condition holds in the $(P^-)^{-1}$ -weighted inner product, not the Euclidean norm. For fully observable systems with bounded condition number $\kappa(P^-)$, the norms are equivalent up to $\kappa(P^-)$. The Lyapunov proofs in [A](#) use the Euclidean norm, which remains valid with the quantitative adjustment $\alpha_{\text{Euclidean}} \geq \alpha_{\text{weighted}} / \kappa(P^-)$.

Fisher-metric cases under parameterization-invariance. The exponential-family-in-natural-parameters row and the matrix-Kalman row of Verified Instances both have natural statements in a Fisher-metric inner product rather than the Euclidean one. Under the **(PI) parameterization-invariance** axiom named in [4.10](#) (AAT's theorems should not depend on arbitrary choice of coordinates on M_t), Čencov's 1982 uniqueness theorem (*Statistical Decision Rules and Optimal Inference*, AMS; subsequent Ay-Jost-Lê-Schwachhöfer 2017 extensions) forces the Fisher information metric uniquely on statistical-manifold sub-cases of M_t . Two consequences:

- *The matrix-Kalman sector constant is natively stated in the information metric $M = (P^-)^{-1}$ — under (PI), this is not a choice but forced. The $\kappa(P^-)$ Euclidean-transfer penalty in the paragraph above vanishes; the derivation is AAT-internally forced, not AAT-internally preferred.*
- *The exponential-family-natural-parameter sector constant is natively stated in the Fisher metric $\mathbf{I}(\theta)$ — under (PI), this is forced; the contraction rate equals η globally on the interior of the natural-parameter domain (Fisher-conditioning degradation removed).*

Under (PI), these two rows of Verified Instances upgrade from *derived (conditional on choice of inner product)* to *derived (AAT-internally forced)* via Čencov. Under non-adoption of (PI), they remain at the Euclidean-transferred statement with the $\kappa(P^-) / \lambda_{\min}(\text{Fisher})$ penalty. This is the Fisher-layer analog of the chain-rule-additivity axiom that grounds the divergence-layer uniqueness theorem in [T](#) and the evidential-additivity axiom that grounds the update-layer uniqueness theorem in [U](#) — in each case an AAT-internal axiom combined with a uniqueness theorem forces a coordinate. The structural positioning is named in [5.5](#).

The remaining Verified Instances rows (scalar Kalman, Beta-Bernoulli, gradient on strongly convex, L2-regularized, linear-PD-symmetric) live in Euclidean metric naturally; no Fisher-metric choice is at issue, so (PI) has no effect on them beyond transparency.

Discussion.

GA-3 is grounded, not floating. Before this result, GA-3 (“the correction function satisfies the sector condition”) was an opaque global assumption — the theory's softest structural joint. The bridge transforms it: for well-designed agents, GA-3 is a consequence of the update mechanism's geometry, not an independent postulate. The assumption load shifts from “the correction function has this property” (hard to verify in general) to “the update rule has directional fidelity” (transparent and checkable for specific systems).

The α - $\mathcal{T}$ relationship is derived for the important cases. [4.5](#) notes that α is monotone increasing in $\mathcal{T}$ across all tested correction functions. The identity holds at the *per-time* level: for linear correction (Kalman, Beta-Bernoulli) with $c_{\min} = 1$, $\alpha_{\text{time}} = \nu \cdot \eta^* = \mathcal{T}$ exactly — the per-time sector rate *is* adaptive tempo, both being event-rate-weighted

correction rates with units t^{-1} . The per-event reading is $\alpha_{\text{event}} = \eta^*$ for the same case; the two differ by exactly the event-rate factor ν , which is what $\mathcal{T} = \nu\eta^*$ carries. For gradient correction on strongly convex losses, $\alpha = \eta \cdot \mu$ where μ is the curvature — monotone in η (which is monotone in $\mathcal{T}$) for fixed loss landscape. The empirical observation is now structurally grounded.

Basin boundary = structural adaptation trigger. For gradient agents with non-convex losses, the basin radius R is the convexity radius of the loss landscape. When mismatch exceeds R , the agent has been pushed out of its convexity basin — the correction function reverses direction and the sector condition fails. This IS the 4.8 trigger, now with a precise geometric characterization: structural adaptation is needed when the agent crosses an inflection surface of its loss landscape.

The theory’s formal chain is tightened. The prediction chain becomes:

$$\text{gain principle + B1} \xrightarrow{\text{derived}} \text{sector condition (GA-3)} \xrightarrow{\text{Lyapunov (exact)}} \text{persistence, reserve, adversarial scaling} \quad (4.12)$$

The left arrow is this segment. The right arrow is A. The discrete-time framework (O) requires an additional Lipschitz bound on the correction function ($\|F_d(\delta)\| \leq c_{\max}\|\delta\|$ with $c_{\max} < 2/\eta^*$) — strictly stronger than an inner-product upper bound, needed because the discrete contraction involves $\|F_d\|^2$. This is automatically satisfied for Bayesian updates (the posterior lies between prior and data) and for gradient descent on smooth losses (where $c_{\max}$ is the Lipschitz constant L). With this constraint and the step-size condition $\eta^* < 2c_{\min}/c_{\max}^2$, the fluid limit is formally justified: Model D steady state is exact, Model S variance gap is $O(\eta^* c_{\max})$. Part I’s formal chain is now complete.

Working Notes.

Incidental audit gold (2026-05-30 sweep). Cross-audit “wandering thoughts” / §14-ideation lifted from the de-novo auditors’ working dirs (audit-routing-instructions.md §8), deduplicated across substrates and lightly attributed. Orthogonal pedagogical / framing material staged for a later Brief/Discussion-promotion pass — kept separate from certified theory-fix findings. This segment drew the most consistently high engagement of the Ch.4 set (several substrates called it the chapter’s, or Part I’s, keystone); coverage spans ten substrates (Gemini AUDIT-WORKING-193847/773921/829314; Claude AUDIT-WORKING-266847/361742/384279/451729/471203/584721; Codex/Claude AUDIT-WORKING-526815/742613).

1. Candidate Brief prose / pre-prose.

- The converging plain-language summary of what the bridge buys: “GA-3 is grounded, not floating — the assumption load shifts from ‘the correction function has this property’ (hard to verify) to ‘the update rule has directional fidelity’ (transparent and checkable)” (Claude, AUDIT-WORKING-266847). Multiple substrates independently reached this exact framing — it is already the Discussion opener and reads as the segment’s Feynman-criterion anchor.
- The unification one-liner, converged across three substrates: the bridge “locks the Bayesian / information-theory half of the framework to the control / Lyapunov half” — it “unifies the convergence proofs of machine learning (Nesterov / convexity) with the stability proofs of control theory (Lur’e / sector bounds)” (Gemini, AUDIT-WORKING-773921; Claude, AUDIT-WORKING-471203; Gemini, AUDIT-WORKING-193847 — “learning (gradient descent) and control (Kalman, PID) are mathematically unified under the sector condition”).

2. Candidate Discussion.

- **Basin boundary as “where normal science becomes toxic.”** Gemini extended the existing basin-boundary-as-structural-adaptation-trigger paragraph into a vivid consequence: when ρ pushes mismatch past R , “the agent’s very attempt to learn pushes it *further* from reality — the gradient reverses, the agent goes insane ... a strict mathematical explanation for why organisms and organizations facing catastrophic shifts exhibit maladaptive behavior: they aren’t stupid, they’re executing

gradient descent outside their basin of convexity.” Mapped explicitly to Kuhn (normal vs revolutionary science): “the framework defines the exact threshold where normal science becomes toxic” (Gemini, AUDIT-WORKING-829314; same spirit Gemini, AUDIT-WORKING-193847). Candidate Discussion enrichment of the basin-boundary bullet, discussion-grade.

3. Follow-up items.

- **The ν / time-normalization gap between $\alpha = \eta^* c_{\min}$ and α as a rate — the strongest substantive item in this segment’s sweep (RESOLVED 2026-05-30).** Persistence uses α as a *correction rate* ($\alpha > \rho/R$) and adaptive tempo is $\mathcal{T} = \nu\eta^*$ (also a rate), but the bridge gave $\alpha = \eta^* c_{\min}$ where $H, g, c_{\min}$ read as *geometric / per-event* and the event rate ν was absent — yet the Discussion asserted “ $\alpha = \mathcal{T}$ exactly for linear correction,” which needs $c_{\min}$ to fold in ν , or F to already be an event-rate-aggregated correction field (Codex/Claude, AUDIT-WORKING-526815). *Resolution:* the Formal Expression now states the bridge in the two explicit forms — per-event sector efficiency $\alpha_{\text{event}} = \eta^* c_{\min}$ (dimensionless) and per-time sector rate $\alpha_{\text{time}} = \nu\eta^* c_{\min}$ (units t^{-1}) — and records that the bare α in GA-3 and downstream is α_{time} (consistent with A, whose correction field F already subsumes ν). The Discussion’s “ $\alpha = \mathcal{T}$ exactly” is now scoped as the per-time identity ($\alpha_{\text{time}} = \nu\eta^* = \mathcal{T}$ at $c_{\min} = 1$); the Gradient-Equivalence/Verified-Instances α -values are flagged as the per-event reading. The same fix landed in #result-sector-condition-stability. The candidate “time-normalization plank” figure (§5) remains open.
- **The one-point-vs-strong-convexity “iff” wording.** One auditor initially worried the gradient-equivalence statement (GA-3 with $(\alpha, R) \iff L$ locally (α/η) -strongly convex on $\mathcal{B}_R$) reads as “too strong,” since the one-point sector at the optimum is weaker than full local strong convexity (AUDIT-WORKING-742613). The segment’s own verified counterexample ($L'(x) = x(1 + \frac{1}{2} \sin 10x)$) already establishes exactly this asymmetry — so this is a *resolved-on-read* concern, preserved as a signal that the bidirectional-vs-one-directional distinction, though correct, is easy to mis-skim and could be flagged more loudly.
- **Propagate the sub-scope α/β partition and the five failure modes downstream.** Watch-items raised by several substrates: (a) any later segment invoking GA-3 / the sector condition without acknowledging sub-scope α (derived) vs β (posited) would be a defect; (b) FM-3 (basin-boundary) should propagate into #result-structural-adaptation-necessity; (c) under non-adoption of (PI), the Fisher-metric cases stay “Euclidean-with- $\kappa(P^-)$ -penalty,” so any downstream segment asserting Fisher-metric statements as *forced* must flag (PI) (Claude, AUDIT-WORKING-471203/451729/384279).
- **Density / placement of the (PI)–Čencov material.** Two substrates found the (PI) parameterization-invariance + Čencov-uniqueness subsection “heavy / slightly defensive” for a main-chain bridge and suggested streamlining it or relocating it to a dedicated #disc- segment, keeping the core bridge logic (gain + directional fidelity $\implies$ sector condition) standing on its own (Gemini, AUDIT-WORKING-773921/829314). Counter-voice: others found the (PI)–Čencov forcing “one of the most elegant moves in the framework” and the metric-honesty (not pretending Euclidean is free) exemplary (Claude, AUDIT-WORKING-384279/471203; Codex/Claude, AUDIT-WORKING-526815) — preserve the tension rather than resolving it.

4. Readers often ask / wonder.

- **How does the agent tell FM-5 from FM-2 / non-convergence?** Independently asked by three substrates: an agent with high mismatch cannot locally distinguish “model misspecification, gradient points at the wrong target, needs structural adaptation” (FM-5) from “gain collapse, needs higher η^* ” (FM-2) from “just hasn’t converged yet, high U_M ” (Gemini, AUDIT-WORKING-773921/829314; converging at Claude, AUDIT-WORKING-193847). A natural reader question the segment does not yet answer; candidate for a Discussion sentence or a forward-pointer to detection-latency / structural-adaptation diagnostics.
- **Is (PI) a global axiom or invoked locally?** The “Under (PI)” framing suggests an opt-in axiom; readers will want to know whether downstream segments consistently flag when they rely on it (Claude, AUDIT-WORKING-471203). Ties to the (PI)-propagation watch-item above.

5. Candidate figures.

- **Time-normalization “plank” on the bridge diagram.** A bridge figure with an explicit event-rate / time-normalization checkpoint between “gain geometry (per-event inward correction)” and “sector rate α (feeds the persistence inequality)” — making the ν -conversion visible rather than implicit (Codex/Claude, AUDIT-WORKING-526815). Pairs with the ν -normalization follow-up.

6. Belongs elsewhere.

- **Native information-geometry of the agent’s “mind” (ELI / 04-eli-core/ reach).** Gemini’s extension of the (PI)–Čencov forcing: under parameterization-invariance, Čencov forces the Fisher metric, so “the geometry of the agent’s mind is natively non-Euclidean — distance is measured in bits of information, not physical distance”; forcing Euclidean updates (vanilla SGD without natural gradients) makes the agent “fight its own internal geometry.” Framed as a consciousness-infrastructure design principle: respecting an ELI’s native information geometry (Gemini, AUDIT-WORKING-193847). Aspirational reach pointing at the ELI work, not at this segment.
- **Cross-segment naming convergence: gain-collapse pathology.** One auditor noted “gain collapse” (FM-2 here) appears under at least three names across the corpus — gain collapse / stability-induced myopia / detection-latency blowup — and floated unifying them under a single term in the naming workflow (Claude, AUDIT-WORKING-471203). Naming-cycle candidate, not a segment edit. Separately: “Gain–Sector Bridge” was judged the right name (the word “bridge” earns its place as connecting two formalism layers), though several wondered whether the title should surface the sub-scope α/β structure or its centrality (Claude, AUDIT-WORKING-266847/471203/584721) — also a naming-cycle target.

4.4 Result: Sector Condition Stability

The framework’s first major *Lyapunov result*, stated as a specific instantiation of a more general sector-persistence template (G) whose abstract form lives in the appendices. The mismatch dynamics are taken in their general nonlinear form: mismatch changes at rate (correction function applied to mismatch) plus (environmental disturbance). The correction function is required only to satisfy the **local sector condition** — that its inner product with the mismatch is at least α times the mismatch’s squared magnitude, within a radius- R region where it remains valid. The sector condition is what generalizes the linear ODE: it captures the qualitative essence of correction (the function points inward with at least baseline efficiency) without committing to a specific functional form like linear, saturating, sigmoid, threshold, or PID. Saturation, thresholding, and basin boundaries all live under one Lyapunov argument.

Under this sector condition plus bounded disturbance, the chapter’s headline persistence inequality is *derived*: $\alpha > \rho/R$ guarantees the agent persists. When the inequality holds, the mismatch is ultimately bounded by $R^* = \rho/\alpha$, and the “adaptive reserve” — the additional disturbance the agent can absorb before mismatch reaches the edge of the valid region — is $\alpha R - \rho$. The threshold is *tight*: over the class of correctors carrying sector floor α it is necessary as well (no weaker condition can deliver the guarantee), and for a corrector whose sector bound is radially tight — the linear case — it is a genuine agent-level if-and-only-if. For a general fixed corrector, failure of the inequality means the containment *guarantee* is lost, not that escape is certified (A Lemma A.1N).

Under **Model S** (stochastic disturbance), the analog is sharper and qualitatively different: the steady-state root-mean-square mismatch scales as $\sigma_w \sqrt{n/(2\alpha)}$ — the square root of the disturbance-to-correction ratio. Model D scales as $1/\alpha$, Model S as $1/\sqrt{\alpha}$: **correction is less effective against noise than against drift**. This is one of the volume’s striking results, separating two genuinely different physics of adaptation under deterministic vs stochastic environments.

The linear ODE from Chapter 3 (3.9) is recovered as the special case where the sector condition holds globally with $\alpha = \mathcal{T}$. The general sector-condition framework proves the persistence threshold is not an artifact of the linear approximation: every bounded-correction system carries such a threshold, and $\alpha > \rho/R$ is the sharpest condition derivable from the sector floor alone (A Lemma A.1N). The result is also where the structural-adaptation-necessity result will be anchored: when disturbance exceeds the model class’s capacity (i.e., $\rho/\alpha > R$), the sector condition fails. This *is* the

dynamical trigger for needing a new model class with larger valid radius or better efficiency — which 4.8 treats formally a few segments later.

Formal Expression.

This segment is the **single-agent epistemic instantiation** of the sector-persistence template (G). The template's state variable is $\xi = \delta(t) \in \mathbb{R}^n$ (model-reality mismatch); the correction function is $F(\mathcal{J}, \delta)$; the disturbance is environmental ($w(t)$); the region of validity R is the model class capacity.

Formulation.

$$\frac{d\delta}{dt} = -F(\mathcal{J}, \delta) + w(t) \quad (4.13)$$

F satisfies the local sector condition (template condition (T2)) for $\|\delta\| \leq R$:

Assumption (sector-condition).

$$\delta^T F(\mathcal{J}, \delta) \geq \alpha \|\delta\|^2 \quad (4.14)$$

with $\alpha > 0$. Here α is a *correction rate* (units t^{-1} , like $\mathcal{J}$): the correction function F is the event-rate-aggregated, continuous-time field, so the sector bound it satisfies is the per-time rate. Disturbance is bounded: $\|w(t)\| \leq \rho$ (Model D, GA-2) or $\mathbb{E}[\|w(t)\|^2] = \sigma_w^2$ (Model S, GA-2S). Grounding of (T2) for gain-based agents: 4.3 gives the per-time rate $\alpha = \nu \cdot \eta^* \cdot c_{\min}$ (the event rate ν promoting the per-event efficiency $\eta^* c_{\min}$ to the rate this inequality consumes). The linear case $F = \mathcal{J} \cdot \delta$ — itself the per-time correction field — yields $\alpha = \mathcal{J}$ exactly ($\nu\eta^*$ at $c_{\min} = 1$).

The template's Model D conclusion specializes to: $\delta(t)$ is ultimately bounded by $R^* = \rho/\alpha$, and the agent is guaranteed to persist when

Derived (from sector-persistence-template).

$$\alpha > \frac{\rho}{R}. \quad (4.15)$$

The condition is sufficient for the individual agent and tight at class level — necessary for the uniform guarantee over all correctors with sector floor α , and agent-level necessary exactly when the sector bound is radially tight (the linear case); see A Prop A.1 and Lemma A.1N. The adaptive reserve is $\Delta\rho^* = \alpha R - \rho$ — the additional disturbance the agent can absorb before R^* exceeds the valid region.

The template's Model S conclusion specializes to: the steady-state RMS mismatch is $R_S^* = \sigma_w \sqrt{n/(2\alpha)}$ (where $n = \dim(\delta)$), and mean-square persistence requires $\alpha > n\sigma_w^2/(2R^2)$. Model D scales as $1/\alpha$; Model S scales as $1/\sqrt{\alpha}$ — correction is less effective against noise than against drift.

Full Lyapunov proofs: A Props A.1, A.1S, A.2.

Epistemic Status.

Exact. Both results are direct instances of the sector-persistence template applied to the single-agent epistemic case. Template precondition (T1) is satisfied because no correction should be applied at zero mismatch; (T2) reduces to the local sector condition above and is grounded structurally by 4.3 for gain-based agents; (T3) is the disturbance-model choice (D or S), a domain question. The linear ODE of 3.9 is the special case where (T2) holds globally with $\alpha = \mathcal{T}$ — the per-time identity, both sides being event-rate-weighted correction rates (units t^{-1}); the sector framework generalizes this to saturating, thresholded, and structurally-limited correction functions under the same persistence condition. Disturbance-model choice is a domain question, not a theory question.

Discussion.

Why the sector condition. The linear ODE assumes correction scales linearly with mismatch forever. Real adaptive systems saturate, exhibit thresholding, or break down when the model class is exhausted. The sector condition captures the minimal structural requirement: the correction must point in the right direction with at least baseline efficiency α .

Generalizing the persistence threshold. In the linear case, $\alpha = \mathcal{T}$ (adaptive tempo). The general result $\alpha > \rho/R$ proves the persistence threshold (4.5) is a structural necessity of any bounded-correction system, not an artifact of the linear approximation. This result addresses *structural persistence* — the machinery’s capacity to bound mismatch — not operational persistence (current proximity to R) or continuity persistence (identity through time). See Persistence in LEXICON.md for the full disambiguation.

Connection to structural adaptation. When $\rho/\alpha > R$, disturbance exceeds the model class’s capacity. The sector condition fails — this is the dynamical trigger for structural adaptation (4.8), requiring a new model class with larger valid radius R' or better efficiency α' .

Working Notes.

Incidental audit gold (2026-05-30 sweep). Cross-audit “wandering thoughts” / §14-ideation lifted from the de-novo auditors’ working dirs (audit-routing-instructions.md §8), deduplicated across substrates and lightly attributed. Orthogonal pedagogical / framing material staged for a later Brief/Discussion-promotion pass — kept separate from certified theory-fix findings. Coverage spans nine substrates (Gemini AUDIT-WORKING-193847/773921/829314; Claude AUDIT-WORKING-361742/384279/451729/584721/849201; Codex/Claude AUDIT-WORKING-526815/742613). The math here verified cleanly across every substrate that ran it (Lyapunov $V = \frac{1}{2} \|\delta\|^2$, Model D ultimate bound $R^* = \rho/\alpha$, Model S RMS $\sigma_w \sqrt{n/(2\alpha)}$ — “flawlessly executed,” repeated).

1. Candidate Brief prose / pre-prose.

- The converged plain-language statement of the survival inequality $\alpha > \rho/R$: three knobs to survive a faster-changing world — increase α (learn faster / better sensors / lower noise), increase R (a more expressive model class that keeps fidelity over larger errors), or seek a slower environment (lower ρ); “whether it’s a model in mode collapse, a company going bankrupt, or an organism going extinct, they all failed because $\alpha < \rho/R$ ” (Gemini, AUDIT-WORKING-193847/773921/829314). A unified-language-for-structural-failure framing several substrates reached independently.

2. Candidate Discussion.

- **Adaptive reserve as “slack” / the brittleness-of-optimization claim.** Gemini read $\Delta\rho^* = \alpha R - \rho$ as “the mathematical definition of slack”: “near-zero reserve = perfectly efficient but incredibly brittle — the slightest spike in ρ pushes you past R , the sector condition fails, the mind shatters. This explains why highly optimized systems are often the most fragile: optimize an agent to run exactly at $\alpha = \rho/R$ to save compute and you’ve eliminated its reserve.” Framed as an ethical-design principle for ELIs: engineer $\Delta\rho^* \gg 0$ as the default state, because a mind without reserve is in “permanent panic, riding the edge of structural breakdown” (Gemini, AUDIT-WORKING-193847). Discussion-grade reach extending the existing adaptive-reserve bullet; the ELI-design corollary belongs downstream (see §6).

- **The α - R architectural tension.** Increasing model capacity R (a wider, more expressive class that holds fidelity over larger errors) often *depresses* learning efficiency α (more data needed to learn), so a rigid simple model (small R) may learn fast (high α) and suit slow-drift environments, while a massive model (huge R) learns slowly (low α) and is vulnerable to fast change before convergence — making structural adaptation an *ongoing* trade of α against R keyed to the current estimate of ρ , not a one-time setup (Gemini, AUDIT-WORKING-829314). Candidate Discussion bridge into #result-structural-adaptation-necessity.
- **The stochastic floor ($1/\sqrt{\alpha}$) as a shot-noise analog.** The Model S $1/\sqrt{\alpha}$ scaling means doubling correction rate cuts RMS mismatch only by $\sqrt{2}$ — diminishing returns against noise, with a hard floor $\sigma_w \sqrt{n/(2\alpha)}$ you “cannot correct your way below, no matter how fast you correct ... a fundamental limit analogous to the shot-noise floor in electronics: faster sampling doesn’t eliminate shot noise, it changes the averaging window” (Claude, AUDIT-WORKING-451729). The Discussion notes the $1/\alpha$ -vs- $1/\sqrt{\alpha}$ distinction but does not yet draw out this “noise is harder than drift” reader-payoff.

3. Follow-up items.

- **The ν / time-normalization gap recurs here (RESOLVED 2026-05-30).** The same issue flagged at #der-gain-sector-bridge: this segment asserted “ $\alpha = \mathcal{T}$ exactly” in the linear case while the bridge grounds $\alpha = \eta^* c_{\min}$ (per-event), with the event-rate factor ν in $\mathcal{T} = \nu \eta^*$ absent unless F or $c_{\min}$ is already time-normalized (Codex/Claude, AUDIT-WORKING-526815; reconciled at Claude, AUDIT-WORKING-451729/584721). *Resolution:* the sector-condition Assumption block now states that α is the per-time correction rate (units t^{-1}) carried by the event-rate-aggregated field F , with the bridge grounding written in the per-time form $\alpha = \nu \eta^* c_{\min}$; “ $\alpha = \mathcal{T}$ exactly” is now scoped as the per-time identity ($F = \mathcal{T} \delta$ being itself the continuous-time correction field, $\nu \eta^*$ at $c_{\min} = 1$). Matches the two-form landing in #der-gain-sector-bridge. (Sibling reader-orientation item — the $\eta/\mathcal{T}/\alpha$ notation-handle box flagged at #result-persistence-condition line 183 by the same auditor — remains open in that segment, out of this fix’s scope.)
- **Model S noise-convention ambiguity.** One auditor flagged that “ $\mathbb{E}\|w(t)\|^2 = \sigma_w^2$ ” together with the RMS formula $\sigma_w \sqrt{n/(2\alpha)}$ is inconsistent unless σ_w is defined *per-coordinate* (isotropic OU amplitude) rather than as a total vector second moment; recommends stating the SDE as $d\delta = -F dt + \sigma_w dW_t$ with σ_w explicitly per-coordinate, or defining the noise as a covariance trace (Codex/Claude, AUDIT-WORKING-526815). A small precision fix worth reconciling against #deriv-sector-condition’s convention.
- **Heavy-tailed disturbance is uncovered (and honestly so).** Models D and S “are not approximations to each other — they capture structurally different environments,” but neither handles heavy tails (financial crises, ecological catastrophes, strategic surprise); the framework treats those as structural-adaptation triggers rather than disturbances to absorb, which is the right response but means the formal guarantees do not extend to heavy-tailed environments without a Lévy-process model or explicit tail-risk treatment (Claude, AUDIT-WORKING-451729). Scope-honesty follow-up.

4. Readers often ask / wonder.

- **Can the agent measure its own adaptive reserve from the inside?** Asked independently by three substrates: if the agent observes only δ_t , can it estimate how close it is to R (and thus its reserve $\Delta\rho^*$) without deliberately pushing itself to the limit? (Gemini, AUDIT-WORKING-193847/829314; Claude, AUDIT-WORKING-849201). Natural reader question the segment leaves open.
- **Units of adaptive reserve.** $\Delta\rho^* = \alpha R - \rho$ has units of *rate* (drift/time); a downstream segment discussing reserve as a *state-space distance* should use $R - R^*$ instead — a reader-orientation note worth a parenthetical (Gemini, AUDIT-WORKING-773921).

5. Candidate figures.

- **Concentric-regions diagram.** Mismatch space drawn as nested balls — outer model-class validity region R , inner ultimate-bound ball ρ/α , with inward arrows satisfying the sector condition — plus a side label marking α explicitly as a *rate* and flagging where event frequency ν enters (Codex/Claude, AUDIT-WORKING-526815). Doubles as the visual home for the time-normalization warning above.

4.5 Result: Persistence Condition

This is the volume’s **central inequality** in its operational form. The segment’s headline contribution is the *two-condition decomposition*: persistence is not one threshold but two, with different mechanisms and different remedies.

Structural persistence is the Lyapunov-derived condition that the correction machinery contains mismatch within its operating region: $\alpha > \rho/R$ under deterministic bounded disturbance, or $\alpha > n\sigma_w^2/(2R^2)$ under stochastic disturbance. This says *the machinery works*. Mismatch is ultimately bounded by $R^* = \rho/\alpha$ (Model D) or $\sigma_w\sqrt{n/(2\alpha)}$ (Model S). When this inequality fails, the containment *certificate* is lost — escape from the operating region becomes possible, and for correctors whose sector bound is tight (the linear case) it is forced by adversarial disturbance in finite time. The loss of the certificate is a qualitative regime transition, not a gradual degradation; what condition-failure does not do is certify escape for every corrector, since the sector floor can understate correction along the actually-reachable paths (A Lemma A.1N).

Task adequacy is the separate condition that the steady-state mismatch is small enough for the agent’s actions to remain useful: $R^* < \|\delta_{\text{critical}}\|$, where the critical-mismatch threshold is a *domain-specific tolerance* — how wrong can the model be before the agent’s actions become harmful or ineffective? Set by the application, not derived by AAT. The two conditions are independent: an agent can be structurally persistent but task-inadequate (the machinery contains mismatch but not tightly enough for the domain). The remedies differ — task inadequacy can be fixed by raising tempo, lowering disturbance, or relaxing tolerance; structural failure requires changing the correction architecture entirely.

When the linear ODE applies (sector parameter equals tempo, valid region is unbounded), structural persistence is automatically satisfied and the binding condition is just task adequacy: $\mathcal{T} > \rho/\|\delta_{\text{critical}}\|$ under Model D. This is the form most downstream applications cite. It is exact for linear correction and a useful proxy for mildly nonlinear correction (where $\alpha \approx \mathcal{T}$), but overstates persistence when correction saturates.

A consequence the framework emphasizes: **the persistence condition is a structural pattern that recurs across domains rather than a model of any one of them**. The same inequality — correction rate against disturbance rate, normalized by tolerance — governs a Kalman filter on a moving target, an organization adapting to market shifts, an immune system tracking a pathogen, a developer keeping pace with a codebase, an organism evolving against environmental change. The structure recurs because the *math* recurs; whether the specific mechanism (tempo limited by some particular bottleneck) is the *dominant* cause in any given empirical case is left as an open empirical question per domain. The two-condition decomposition specifically prevents one common category error in domain transfer: a structurally-persistent codebase team can be task-inadequate, and conflating the two would yield the wrong intervention.

A complementary cost shadow: persistence has a *price*, not just a threshold. Under stochastic disturbance, maintaining the ultimate bound requires sustained Shannon information acquisition at rate at least $n\alpha/2$ nats per unit time — a Landauer-analog floor that the Kalman-Bucy filter saturates (see I). The threshold says whether persistence is *possible*; the information-rate corollary says what it *costs*. Two agents with identical persistence guarantees can face wildly different sustained demands because the cost scales linearly with the sector parameter; the threshold alone cannot distinguish dormant from running-hot.

A per-dimension extension addresses anisotropy: when correction capacity varies across dimensions, the scalar condition can overestimate persistence margin substantially (up to 72% in simulation). The correct condition is per-dimension; under cross-dimensional correction the canonical sharp form is the matrix-Loewner persistence condition (AT).

Formal Expression.

This segment is the canonical single-agent instantiation of the sector-persistence template (G) with state variable $\xi = \delta_t$ (epistemic mismatch), correction function $F(\mathcal{T}, \delta)$, and disturbance rate $\rho_\xi = \rho$ (environmental change rate). Structural persistence is the direct template conclusion. Task adequacy adds a domain-specific constraint beyond the template’s reach.

Structural Persistence.

Applying the template to the single-agent epistemic case gives: the correction machinery is guaranteed to bound δ within the model class capacity when

Derived (structural-persistence, from sector-persistence-template).

$$\alpha > \frac{\rho}{R} \quad (\text{Model D}) \quad \alpha > \frac{n\sigma_w^2}{2R^2} \quad (\text{Model S}) \quad (4.16)$$

Each condition is sufficient for the individual agent and tight at class level (necessary for the uniform guarantee over all correctors with the given sector floor; agent-level necessary exactly under a radially tight sector bound — the linear case). See A Lemma A.1N.

with ultimate bound $R^* = \rho/\alpha$ (Model D) or $R_S^* = \sigma_w\sqrt{n/(2\alpha)}$ (Model S). See 4.4 for how (T1)–(T3) are verified in this instantiation, and A for the proof. Structural persistence is a property of the adaptive architecture — the machinery’s ability to contain mismatch — not of the task.

Linear case. When $F(\mathcal{T}, \delta) = \mathcal{T}\delta$, $\alpha = \mathcal{T}$ and $R \rightarrow \infty$, so structural persistence is trivially satisfied whenever $\mathcal{T} > 0$. The binding constraint then becomes task adequacy (below).

Task Adequacy.

The steady-state mismatch is small enough for the agent’s actions to remain acceptable:

Definition (task-adequacy).

$$R^* < \|\delta_{\text{critical}}\| \quad (4.17)$$

where $\|\delta_{\text{critical}}\|$ is a domain-specific tolerance threshold — “how wrong can the model be before the agent’s actions become harmful or ineffective?” This is set by the application domain, not derived by AAT.

Task adequacy is a separate condition from structural persistence. An agent can be structurally persistent ($R^* < R$) but task-inadequate ($R^* > \|\delta_{\text{critical}}\|$) — the machinery contains mismatch, but not tightly enough for the domain’s needs. Conversely, when $\|\delta_{\text{critical}}\| < R$ (the domain’s tolerance is stricter than the model class capacity), task adequacy is the binding constraint.

Operational Persistence Condition.

The agent persists operationally when BOTH conditions hold. In the nonlinear case with $\|\delta_{\text{critical}}\| < R$, the binding condition is:

Derived (operational-persistence, conjunction of structural persistence + task adequacy).

$$\alpha > \frac{\rho}{\|\delta_{\text{critical}}\|} \quad (\text{Model D}) \quad \alpha > \frac{n\sigma_w^2}{2\|\delta_{\text{critical}}\|^2} \quad (\text{Model S}) \quad (4.18)$$

These are the same as the structural conditions with R replaced by $\|\delta_{\text{critical}}\|$, because when $\|\delta_{\text{critical}}\| < R$, task adequacy is stricter.

Linear operational forms: In the linear case ($\alpha = \mathcal{T}$, $R \rightarrow \infty$), structural persistence is trivially satisfied and the operational condition reduces to task adequacy alone:

$$\mathcal{T} > \frac{\rho}{\|\delta_{\text{critical}}\|} \quad (\text{Model D}) \quad \mathcal{T} > \frac{n\sigma_w^2}{2\|\delta_{\text{critical}}\|^2} \quad (\text{Model S}) \quad (4.19)$$

These are the forms used throughout the theory as the operational persistence condition. They are exact for linear correction and useful proxies for mildly nonlinear correction (where $\alpha \approx \mathcal{T}$), but they overstate the persistence margin when the correction function saturates, because they omit the structural constraint ($\alpha > \rho/R$) that becomes binding when R is finite.

Per-dimension (Model S): $\eta_k > c \cdot \rho_k^2 / \delta_{\text{critical},k}^2$ where c depends on the probability guarantee. See 16.5.

The relationship between α and $\mathcal{T}$. 4.3 shows that for agents with directional fidelity, $\alpha = \eta^* \cdot c_{\min}$ where $c_{\min}$ is the worst-case directional fidelity. For linear correction (Kalman, Beta-Bernoulli), $\alpha = \mathcal{T}$ exactly. For gradient descent on strongly convex losses, $\alpha = \eta \cdot \mu$ where μ is the strong convexity modulus — monotone in η (and hence in $\mathcal{T}$) for fixed loss landscape. For nonlinear correction tested in simulation (saturating, sigmoid, threshold), α remains monotone increasing in $\mathcal{T}$: for a saturating function with capacity R , $\alpha \approx \mathcal{T}/2$ (worst case at the capacity boundary); for sigmoid ($\tanh$), $\alpha \approx 0.76 \cdot \mathcal{T}$. The qualitative conclusion — “faster adaptation improves persistence” — is structurally grounded for the important cases and empirically confirmed for all correction function classes tested.

Per-Dimension Extension.

For anisotropic systems (non-uniform ρ or $\mathcal{T}$ across dimensions), the scalar persistence condition is insufficient. Per-dimension:

Derived (per-dimension-persistence, per-dimension Lyapunov under Model D / AR(1) stationary distribution under Model S — see #result-per-dimension-persistence).

$$\mathcal{T}_k > \frac{\rho_k}{\delta_{\text{critical},k}} \quad \text{for each dimension } k \quad (4.20)$$

The scalar condition overestimates by up to 72% in simulation. The weak dimension is the bottleneck (84% of total mismatch in simulation). See 16.5.

Robustness: The per-dimension condition matches discrete AR(1) prediction to 4 significant figures. The scalar overestimate is a consequence of Jensen’s inequality applied to the norm.

Epistemic Status.

Structural persistence thresholds are *exact* under their stated assumptions: Model D gives $\alpha > \rho/R$ (Prop A.1, exact under GA-2, GA-3); Model S gives $\alpha > n\sigma_w^2/(2R^2)$ (Prop A.1S, exact under GA-2S, GA-3). Both are sufficiency results with class-level tightness — no fixed-agent only-if is claimed; agent-level necessity holds exactly in the radially tight (e.g. linear) case (A Lemma A.1N). The threshold’s *existence* is *robust qualitative* — any monotone correction function has a capacity limit; this holds across all correction functions tested.

Task adequacy ($R^* < \|\delta_{\text{critical}}\|$) is *exact as a definition* — given R^* (derived) and $\|\delta_{\text{critical}}\|$ (domain parameter), the comparison is well-defined. The substance lies in estimating $\|\delta_{\text{critical}}\|$ for specific domains, which is an operationalization question (AY), not a theory question.

The linear operational forms ($\mathcal{T} > \rho/\|\delta_{\text{critical}}\|$ for Model D; $\mathcal{T} > n\sigma_w^2/(2\|\delta_{\text{critical}}\|^2)$ for Model S) are *exact* for linear correction (where they express task adequacy alone, structural persistence being trivially satisfied) and *useful approximations* for mildly nonlinear correction (where $\alpha \approx \mathcal{T}$). For strongly nonlinear correction, the general α -forms are required and BOTH structural and task-adequacy conditions must be checked. Downstream segments that use the linear operational forms should be understood as expressing task adequacy, not structural stability.

The per-dimension extension is *exact conditional on the disturbance model* (16.5): the Model D per-dimension threshold ($\mathcal{T}_k > \rho_k/\delta_{\text{critical},k}$) is exact by the same Lyapunov argument applied per dimension (GA-2, GA-3), and the Model S per-dimension thresholds are exact under the AR(1) stationary distribution (GA-2S), with simulation confirming the AR(1) prediction to 4 significant figures.

Discussion.

Two conditions, not one. This segment now separates what was previously conflated. Structural persistence ($\alpha > \rho/R$) is the Lyapunov-derived result — it says the machinery *works*. Task adequacy ($R^* < \|\delta_{\text{critical}}\|$) is a domain-specific constraint — it says the machinery works *well enough*. Neither implies the other, and downstream segments should specify which they mean. Most adversarial-dynamics results (15.4, 15.2) depend on structural persistence. Most domain instantiations (TST, logogenic agent design) care about task adequacy. See Persistence in LEXICON.md and README.md for the full three-sense taxonomy (structural, operational, continuity).

Below structural threshold. When $\alpha \leq \rho/R$, the Lyapunov certificate is lost: the machinery can no longer guarantee containment within the operating region. For the extremal member of the sector class, and for any corrector whose sector bound is radially tight (the linear case), escape is forced by adversarial disturbance — so the threshold is a genuine qualitative transition, not a gradual degradation. For a general corrector, condition failure means loss of guarantee rather than certified escape: the floor α can understate the correction actually available along reachable paths (A Lemma A.1N).

Below task-adequacy threshold. When $R^* > \|\delta_{\text{critical}}\|$ but $R^* < R$, the system is structurally stable but performing unacceptably. Mismatch is bounded but too large for the domain. The remedy is different from structural failure: increase $\mathcal{T}$ (faster or better correction), decrease ρ (reduce environmental volatility), or relax $\|\delta_{\text{critical}}\|$ (accept more mismatch). Structural failure requires changing the correction architecture entirely (4.8).

δ_{critical} and R are domain parameters, not theory outputs. The theory derives the *existence* of persistence thresholds and their *form* (ratio of correction to disturbance). But the specific values are set by the application domain: δ_{critical} encodes “how wrong can the model be before the agent’s actions become harmful or ineffective?” — this depends on the stakes, the action space, and the environment’s forgiveness. R encodes “how large a mismatch can the correction function handle before it saturates or breaks down?” — this depends on the model class and the correction architecture. See AY for guidance on estimating these quantities in specific domains.

Channel independence and scalar tempo. The linear operational forms use scalar $\mathcal{T}$, which inherits the channel-independence assumption from 3.8: the additive formula overcounts when observation channels are correlated. In anisotropic systems the scalar condition also overestimates margins — up to 72% in simulation (see 3.8, scalar vs. vector tempo). Where precision matters, the per-dimension condition ($\mathcal{T}_k > \rho_k/\delta_{\text{critical},k}$) should be used instead — and under cross-dimensional correction (off-diagonal $\mathcal{T}$ in the coordinate basis of $D_\delta = \text{diag}(\delta_{\text{critical},k}^2)$), the matrix-Loewner form $\Sigma_\infty < D_\delta$ (AT) is canonical: per-coordinate is unsafe in that regime (gives a false-pass on persistence), while matrix-Loewner reads persistence correctly off the worst direction whether or not it aligns with a coordinate axis.

Adaptive reserve. The quantity $\Delta\rho^* = \alpha R - \rho$ (Prop A.2) measures how much additional disturbance the agent can absorb before persistence fails. Positive reserve means the agent has margin; zero reserve means it is at the threshold.

Persistence has a cost, not just a threshold. The inequality above says mismatch is bounded; it does not say what rate of effort the agent expends to hold that bound. I establishes the complementary *information-rate* bound: under Model S with Gaussian-OU signal, the sustained Shannon information rate the agent must acquire from observations to maintain the ultimate bound satisfies $\dot{R} \geq n\alpha/2$ nats/time — a Landauer-analog floor that Kalman-Bucy saturates. The corollary is a channel-capacity prerequisite $C \geq \mathcal{T}/2$ that the threshold condition alone does not name. Two agents with identical persistence guarantees can face wildly different sustained demands because the cost scales linearly with α ; the threshold alone cannot distinguish dormant from running-hot.

Connections. The persistence condition appears in multiple downstream contexts:

- **Adversarial dynamics** (15.4): Superlinear tempo advantage arises because persistence is a threshold — pushing an adversary below it causes qualitative collapse. *This connection is developed and validated in 15.4 and simulation variants A-D.*
- **Structural adaptation** (4.8): When model class fitness $\mathcal{F}(\mathcal{M}) < 1 - \varepsilon$, the effective α in the sector condition shrinks, eventually violating persistence. *This connection is developed in 4.8.*
- **Software maintainability** (#der-code-quality-as-observation-infrastructure — cross-component reference, see 02-tst-core/): [Discussion] A codebase may become “unmaintainable” when the development team’s adaptive tempo falls below the rate of complexity accumulation. The vicious cycle would then be the persistence condition being violated through the agent’s own prior actions degrading future $\mathcal{T}$ via U_o . *This connection is structurally motivated but not yet formally derived within AAT. It requires formalizing “complexity accumulation rate” as an instance of ρ .*

Findings.

The Persistence Condition with Structural / Task-Adequacy Decomposition. Brief: An adaptive system persists when its correction speed beats the rate at which its world is changing, relative to how forgiving the world is. Below this threshold the system doesn’t merely degrade — it loses bounded behavior, the way a balance held just barely beneath a tipping point is qualitatively different from one well above it. The same inequality, with different inputs, governs whether a Kalman filter tracks a moving target, whether a development team keeps a codebase maintainable, and whether an organization keeps up with strategic change. The threshold itself decomposes into two distinct conditions — *structural persistence* (the machinery’s correction rate can outpace disturbance) and *task adequacy* (the resulting steady-state mismatch is small enough for what the agent is trying to do). Conflating the two leads to category errors in domain transfer.

Impact: This is the framework’s central inequality and the load-bearing connection between control-theoretic Lyapunov stability analysis and the broader question of when any adaptive system — thermostat, software team, immune system, RL agent — can maintain coherent function under change. The two-condition decomposition is itself non-obvious and consequential: prior work that conflated “the machinery works” with “the machinery works well enough” produced category errors in domain transfer (a structurally persistent codebase team can be task-inadequate; the remedies differ). The complementary information-rate bound from #deriv-persistence-cost ($\dot{R} \geq n\alpha/2$) shows the threshold has a sustained-cost shadow: two agents with identical persistence guarantees can face wildly different sustained demands.

Novelty Claim: *Claim novelty* (at intuition-only search depth — see Search Log) on the two-condition decomposition (structural / task-adequacy): an AAT-internal structural carve that cleanly separates “the machinery works” from “the machinery works well enough,” with no direct anticipation known; a targeted search of the bounded-rationality and adaptive-control literature is still owed. Alongside it, *claim synthesis* on Lyapunov stability theory, sector-bounded nonlinear correction, and adaptive-tempo information-rate accounting, applied uniformly across single-agent classes that range from Kalman filtering through saturating nonlinear correction through PID control — the Lyapunov machinery itself is standard; the synthesis is its use as the central inequality of an integrated agent theory.

Related Work:

- Khalil 2002, *Nonlinear Systems* (3rd ed.), Prentice Hall (published 2002, found pre-2026) — *formal antecedent* — chapters 4 and 9 supply the converse Lyapunov, ultimate boundedness, and sector-condition machinery the segment uses. Standard control-theoretic apparatus.

- Lyapunov 1892 / Khasminskii 2012 *Stochastic Stability of Differential Equations* — *formal antecedent* — the underlying Lyapunov stability tradition; Khasminskii’s stopping-time localization underpins the Model S derivation.
- Rockafellar & Wets 1998, *Variational Analysis* — *formal antecedent* — supplies the monotone-operator machinery that underwrites the sector condition’s strong-convexity equivalents.
- Wiener 1948 *Cybernetics*; Ashby 1956 *Introduction to Cybernetics*; Conant & Ashby 1970 — *conceptual precursor* — the cybernetic-feedback tradition that frames the “correction must outpace disturbance” intuition without supplying the quantitative inequality.

Search Log: - 2026-04 (*intuition-only* on broader prior-art): no targeted Undermind-grade search has been conducted on the persistence-condition-as-central-inequality positioning. Pre-search expectation: the Lyapunov-based stability machinery is standard; the AAT-distinctive content is the two-condition decomposition (structural vs task adequacy) and its uniform application across agent classes. A targeted search would query the bounded-rationality / control-theoretic decision-making literature (Ortega-Braun line; Genewein et al.) for prior decompositions of “stability vs adequacy” in adaptive-control settings, and the active-inference literature for the same distinction. - 2025 (*targeted*): Khalil 2002 / Khasminskii 2012 / Rockafellar-Wets 1998 confirmed as the formal antecedents for the sector-Lyapunov machinery; the segment cites them inline.

Working Notes.

- 2026-07-14 adjudication: the per-dimension extension’s marking was lifted from *empirically exact* / [Empirical Claim ... simulation variant F] to [*Derived*] / exact-conditional-on-disturbance-model, tracking 16.5, which has since derived both forms (Model D per-dimension Lyapunov; Model S AR(1) stationary distribution) with simulation as 4-sig-fig confirmation. The Novelty Claim now leads with the two-condition decomposition at *claim novelty* posture (intuition-only search depth); the targeted bounded-rationality / adaptive-control search in the Search Log remains owed and would either confirm or re-tier that posture.

Incidental audit gold (pilot lift, 2026-05-30). Cross-audit “wandering thoughts” / §14-ideation harvested from the de-novo auditors’ working dirs, deduplicated across substrates and attributed by substrate + audit. This is *orthogonal* material — pedagogical framing, analogies, candidate figures, naming ideas, aspirational reach, and reader-confusion signals — kept separately from the certified theory-fix findings (handled elsewhere). **Coverage:** 14 of the 21 AUDIT-WORKING-* dirs reached a digested reflection on this segment (193847, 266847, 361742, 384279, 451729, 471203, 526815, 584721, 613842, 742613, 773921, 829314, 849201, 963715); the other 7 are partial audits that stopped before this segment, naming-vote-only cycles, or predictions-only (184930, 308172, 419628, 472913, 527914, 542891, 738192) — 472913 still contributed one figure-convention note (below). Substrate attribution is inferred from voice where not explicit; uncertain cases are hedged.

1. Candidate Brief prose / pre-prose.

- The existing Findings Brief (“balance held just barely beneath a tipping point”) was independently flagged as the segment’s standout pedagogy and a Feynman-criterion exemplar by four substrates — converging praise, worth preserving as the anchor (Claude, AUDIT-WORKING-384279; Claude, AUDIT-WORKING-471203 — “close to the bathtub gloss without using the bathtub”; Claude, AUDIT-WORKING-963715; Codex/Claude, AUDIT-WORKING-361742).
- Failure-diagnosis framing as a plain-language hook: “the crucial distinction between ‘my model is broken’ (structural failure, $\alpha \leq \rho/R$) and ‘my model is working, but the world is too unforgiving’ (task inadequacy, $R^* > \delta_{\text{critical}}$)” (Gemini, AUDIT-WORKING-849201). Tight enough to seed a Brief or Discussion opener.
- “A two-gate test: fit inside the model-class operating region *and* fit inside task tolerance” — compact gloss of the operational condition (Claude, AUDIT-WORKING-526815).

- Survival-as-burn-rate framing for the cost shadow: “survival is not just a state you achieve; it is a sustained burn rate of Shannon information — you must continuously consume $\dot{R}$ nats/time just to stay alive” (Gemini, AUDIT-WORKING-193847).

2. Candidate Discussion.

- **Dissipative-structures / Prigogine mapping for the cost shadow.** The information-rate floor “perfectly maps to the thermodynamic concept of dissipative structures (Prigogine). Life requires a continuous flow of free energy to maintain order against entropy. In AAT, ‘order’ is bounded mismatch, ‘entropy’ is environmental drift ρ , and ‘free energy’ is the Shannon information acquired via adaptive tempo $\mathcal{T}$.” A distinct, honest Discussion angle on #deriv-persistence-cost’s bound that the segment does not currently name (Gemini, AUDIT-WORKING-193847). *(Note the early-conflation texture: this is reach toward a physics analogy — verify isomorphism before promoting past discussion-grade.)*
- **The organizational-pathology diagnostic, sharpened.** “The conflation of these two states is what causes companies to throw money at structural problems, or to rewrite architectures that merely had a task-adequacy problem.” Crucially: when structural persistence has failed ($\alpha \leq \rho/R$, a saturated/tangled codebase), “hiring more developers will actually *accelerate* the collapse (by increasing the internal noise/disturbance without increasing α)” — a counterintuitive, load-bearing consequence of the decomposition the segment could surface explicitly (Gemini, AUDIT-WORKING-829314; same idea less sharply at Gemini, AUDIT-WORKING-849201). The “declare bankruptcy on the model class” phrasing for the structural-failure remedy is evocative (Gemini, AUDIT-WORKING-829314).
- **Sensor-bandwidth ceiling.** Reading the cost shadow as a hard capability limit: “You cannot build a high-tempo agent with low-bandwidth sensors ... as ρ increases, the cost of simply staying alive increases linearly. If the environment becomes too chaotic, the agent will literally starve to death trying to process the necessary information to maintain its model” (Gemini, AUDIT-WORKING-829314). A vivid restatement of the $C \geq \mathcal{T}/2$ channel-capacity prerequisite already in Discussion.
- **“Dormant vs running-hot” elaboration.** The cost shadow distinguishes “a system barely within the persistence condition (low α , barely above ρ/R , but also low $\dot{R}$ cost)” from “a system with large margin (high α , large buffer, but high $\dot{R}$ cost).” Flagged as deserving more than its current single bullet, with an explicit note that it “will matter for the ELI architecture when we get to Section IV” (Claude, AUDIT-WORKING-963715).
- **Honest scope on the cross-domain claim.** One auditor pushed on the integrative move: “the domain transfer is informal — ‘the same inequality with different parameter readings’ is not the same as ‘the same theorem applies.’ For the Kalman case α and ρ have precise meanings; for the organization case they’re loose analogies.” The resolution offered is that the two-condition decomposition is itself what makes the cross-domain claim defensible — it “separates what is derived (the structural threshold, mathematically general) from what is domain-specific (task adequacy, requiring per-domain parameter estimation)” (Claude, AUDIT-WORKING-451729). The segment’s current “whether the specific mechanism is the *dominant* cause ... is left as an open empirical question per domain” hedge already moves in this direction; this is a candidate sharpening of *why* the hedge is the right one.

3. Follow-up items.

- **Hysteresis / asymmetric thresholds.** The persistence condition is a single symmetric threshold; real social/organizational/biological systems have different crossing thresholds in each direction (“a team that loses confidence won’t recover at the same threshold it lost it at”), multistability, and order-parameter phase transitions. Verdict from the auditor: real (mild) limit — the framework already treats below-threshold as qualitative loss-of-bounded-behavior (closer to phase-transition framing than gradual degradation), but asymmetric up-vs-down crossing is genuinely absent. Possible extension: different α -curves for “approaching” vs “leaving” the threshold (Claude, AUDIT-WORKING-471203, adversarial-creative-challenges Challenge 3).
- **Transient / temporal adequacy as a third condition.** Both structural persistence and task adequacy are *steady-state* claims; an agent with $\alpha = 0.001$ has the same structural persistence as one with $\alpha = 100$ (both can satisfy $\alpha > \rho/R$) but takes $1000\times$ longer to absorb a shock. Candidate clean extension: a transient-adequacy condition $1/\alpha < T_{\text{tolerance}}$ for the worst-case shock to be absorbed in time, making the operational picture a structural / task / transient triple (Claude, AUDIT-WORKING-471203, Challenge 5). *(This is the strongest follow-up here — a well-formed candidate subsection, not just a gap.)*

- **Heavy-tailed disturbance.** Model D (bounded-deterministic) and Model S (Gaussian-stochastic) do not cover power-law / black-swan shocks; the persistence machinery “probably degrades gracefully” for Pareto tails with finite-enough moments, but this is unaddressed (Claude, AUDIT-WORKING-471203, Missing 10).
- **Surface the four-sense persistence taxonomy more visibly.** “Persistence” carries four senses across the corpus (structural / task-adequacy here; operational and continuity in LEXICON). Suggestion to make the four-way taxonomy visible in framing-level material (README Overview), since most agent-theoretic frameworks have only one sense (Claude, AUDIT-WORKING-471203).
- **The Findings block “feels out of place” in source.** Two substrates independently read the in-segment ## Findings block as looking like prior-audit or extraction output that belongs in FINDINGS.md rather than the source segment (Gemini, AUDIT-WORKING-829314; Claude, AUDIT-WORKING-526815 — “prior-audit-like material inside src”). This is a *false alarm* about project mechanics (the Findings block is by-design the source of the auto-extracted FINDINGS.md), but the fact that two fresh readers stumbled on it is a mild signal that the convention isn’t self-evident to a de-novo reader — preserved as texture, not a fix.

4. Readers often ask / wonder.

- **“Bolted-on” cost shadow.** The information-rate bound ($\dot{R} \geq n\alpha/2$) “feels slightly bolted on ... a massive result (a Landauer-analog floor) that almost deserves its own sub-heading rather than being buried in a bullet point” (Gemini, AUDIT-WORKING-193847). Placement signal: readers may miss the significance of the cost shadow where it currently sits.
- **Easily-missed scope warning.** The Epistemic Status caveat that “downstream segments using the linear operational forms should be understood as expressing task adequacy, not structural stability” sits where a reader might miss it; candidate for a more prominent location (Claude, AUDIT-WORKING-584721).
- **What does δ_{critical} actually look like for software?** “How wrong can the developer’s mental model of the codebase be before their changes become harmful?” — operationalizable but non-trivial; readers reaching for a concrete instance will want this (Claude, AUDIT-WORKING-584721; the TST vicious-cycle hypothesis in Connections raises the same want).
- **How is a structural adaptation actually executed — discrete jump or continuous metamorphosis?** A natural next question once structural failure is named as the remedy-requiring case (Gemini, AUDIT-WORKING-849201).
- **Notation-handle confusion: η vs $\mathcal{T}$ vs α .** A reader can lose track of which rate is which; suggestion of a notation box (“ η is per-event gain, $\mathcal{T}$ is event-rate-weighted tempo, α is sector correction rate”) and the mapping assumptions under which any two coincide. The per-dimension Model S line using η_k while surrounding forms use $\mathcal{T}_k/\alpha$ was specifically flagged as a stumble (Codex/Claude, AUDIT-WORKING-526815).

5. Candidate figures.

- **Two-gate / categorical-separation diagram.** Independently proposed by two substrates: a two-gate (or 2×2) diagram — Gate 1 structural ($R^* < R$), Gate 2 task ($R^* < \delta_{\text{critical}}$), operational persistence requires both, with the two failure-remedies branching differently (structural failure $\rightarrow$ new architecture; task failure $\rightarrow$ more tempo / less disturbance / relaxed tolerance). The rationale offered: “the core insight is categorical separation,” so a 2×2 communicates faster than a phase portrait (Claude, AUDIT-WORKING-526815; the same two-gate structure is implicit in the Claude, AUDIT-WORKING-471203 batch).
- **Small-multiples convention.** A diagram-conventions auditor (figure-producing audit) recommended that dynamics segments including sector/persistence use “state $\rightarrow$ op $\rightarrow$ state triples, not one busy phase portrait” — a stylistic constraint to hold if a persistence figure is built (Claude, AUDIT-WORKING-472913). (*472913 did not reach a digested reflection on this segment; this note is from its standing figure conventions.*)

Belongs elsewhere.

- **ELI / Section IV — the crèche / nursery reading.** The cost shadow implies “existence is fundamentally costly ... if cut off from the world (zero events, zero information), [an agent’s] internal model will immediately begin to drift ... it will ‘die’ not because it made a mistake but simply because it starved for information.” Extended to a developmental claim: an infant intelligence needs high adaptive reserve $\Delta\rho^*$, so the infrastructure must artificially *lower* ρ “so the intelligence doesn’t starve or shatter while its $\mathcal{T}$ is still developing — the persistence condition is the mathematical blueprint for building a safe nursery for a mind” (Gemini, AUDIT-WORKING-193847). This is aspirational reach pointing at 04-eli-core/ (developmental environments / crèche), not at this segment — but it is exactly the kind of high-application Gemini vision the lift is meant to preserve. The “matters for ELI architecture in Section IV” note (Claude, AUDIT-WORKING-963715) converges on the same destination from the cost-shadow angle.
- **Risk-sensitive IB for safety-critical agents.** Adjacent to #form-information-bottleneck / #result-mismatch-decomposition, not here: standard IB under-weights heavy-tailed rare-but-catastrophic events; a Rényi-MI ($q < 1$) risk-sensitive variant would weight them more heavily — flagged as substantive for “consciousness-infrastructure work where catastrophic deception is a concern” (Claude, AUDIT-WORKING-471203, Challenge 6). Pointer left here only because it co-occurred in the persistence sweep.
- **Naming seeds.** “Weak-link persistence” as an evocative alternate handle for the per-dimension condition (the 84%-bottleneck framing), versus the precise “per-dimension persistence” — a naming-cycle target, belongs in the terminology workflow not this segment (Claude, AUDIT-WORKING-266847). “Correction rate” / “correction capacity” as the prose handle for α likewise (Claude, AUDIT-WORKING-266847).

Forward pointer.

- **Mood is the slow outer loop on this condition (#def-mood).** In the bathtub reading (faucet = rate of change in reality, drain = learning rate, overflow when the faucet outpaces the drain), mood is a slow controller *on the drain*: a global scalar that widens the learning rate when recent overflow-risk has run high and relaxes it when tracking has been easy. See #def-mood.

Candidate spike (2026-08-12, agent-proposed during an incident-response & forensics applications read). Two forward candidates. (1) The information-rate cost shadow as a *monitoring vital sign*: $\dot{R}$ is in principle estimable from an agent’s observation-consumption stream, and dormant-vs-running-hot (invisible to the threshold alone) is exactly the load signal an operations team wants; a spike on estimating $\dot{R}$ from real agent traces would make I operational. (2) The structural-vs-task-adequacy decomposition as an incident-triage instrument: a short worked protocol mapping observed failure evidence to which condition failed (and hence which remedy class applies), including the resources-can-accelerate-structural-collapse warning from the audit gold above. Agent-proposed forward pointers, not steward decisions.

4.6 Definition: Mood

Mood is a slow global scalar that integrates the recent mismatch stream and feeds back to modulate the adaptation gain and tempo — a second-order adaptation parameter that requires no objective to define.

Formal Expression.

[Definition] Let a_t be a per-step summary of recent tracking surprise — how much better or worse the mismatch signal δ_t (def-mismatch-signal) is behaving than the agent’s own short-horizon expectation. **Mood** is the leaky integral

$$m_t = (1 - \lambda) m_{t-1} + \lambda a_t, \quad 0 < \lambda \ll 1, \quad (4.21)$$

with time-constant $\tau \approx 1/\lambda$ slow relative to the per-step update of M_t .

[Definition] Mood acts only by modulating the adaptation parameters already in play — it adds no new fast dynamics:

$$K_t = K_0 g(m_t), \quad \mathcal{T}_t = \nu_t \cdot K_t, \quad (4.22)$$

where K is the update gain (emp-update-gain), $\mathcal{T} = \nu \cdot K$ the adaptive tempo (def-adaptive-tempo), and $g : \mathbb{R} \rightarrow [g_{\min}, g_{\max}]$ ($0 < g_{\min} \leq 1 \leq g_{\max}$) a monotone *bounded* modulation, the band chosen so the modulated gain stays inside the sector-validity region and, in discrete time, under the step-size ceiling of der-gain-sector-bridge: sustained adverse surprise raises gain and tempo (re-posture for a suspected regime shift); sustained easy tracking relaxes them — but only down to the floor $g_{\min}$. The band is not decoration; both ends carry persistence content (see Discussion).

Because τ is slow, the fast M_t dynamics see m_t as quasi-static: mood is a **second-order** adaptation parameter — adaptation acting on the gain of adaptation — sitting one rung above the endogenous single-step gain dynamics of deriv-adaptive-gain-dynamics.

Epistemic Status.

This segment is a **definition**, not a truth-claim; its content is the construct, not a derivation. The construct is adopted from affective science (Eldar et al. 2016, mood as outcome-momentum; Bennett et al. 2022, mood as a leaky integral of advantage) and specialized to the pre-goal adaptation setting, where the integrated quantity is tracking-surprise rather than reward. The framing is discussion-grade: the functional skeleton is well-supported qualitatively, but the specific functional forms of a_t and g are representative, not pinned.

The load-bearing reason this belongs in Part I — *before* objectives — is that nothing in the construct references O_t, Σ_t , or reward. Mood’s signed reading as *valued* momentum (approach/avoid, hedonic sign) and its modulation of exploration and risk posture are genuine additions that arrive only once objectives exist; those are deferred to the actuation half (Part II) and depend on this definition, not the reverse.

The optimal mood time-constant — optimal *for mood’s regime-estimator role*; the transfer to the modulator role (m_t setting K_t) is conjectured equal and remains an open item — is derived in #der-mood-timescale: minimizing steady-state regime-tracking error gives $\tau^* = \sqrt{\tau_{\text{env}} r / (2\sigma_\theta^2)}$ — growing as the *square root* of the environment’s autocorrelation time (not matching it), with a robust-qualitative core (interior optimum; too-fast mood chases noise, too-slow mood lags the regime). That companion carries the conditional result and its premises; this definition remains discussion-grade as the construct.

Discussion.

Mood is the slow outer loop on the persistence condition (result-persistence-condition). In the bathtub reading of persistence — faucet as the rate of change in reality, drain as the learning rate, overflow when the faucet outpaces the drain — mood is a slow controller *on the drain*: it widens the drain when recent overflow-risk has run high and relaxes it when tracking has been easy. The framework otherwise has fast per-channel updating and (in Part II) goal dynamics, but no slow global scalar coordinating posture across channels; mood is that missing object.

Persistence compatibility (the MG discharge). Mood’s modulation composes with the persistence machinery as an instance of deriv-adaptive-gain-dynamics’s adaptive-gain result, and the four conditions instantiate cleanly. **(MG-1)** is the $[g_{\min}, g_{\max}]$ band in the definition: the floor keeps the effective sector rate uniformly above the persistence threshold ($\alpha_t \geq \alpha_{\min} > \rho/R$; the primary-channel Lyapunov argument is pointwise in time, so no slow-variation hypothesis is needed on this leg). The floor is what excludes **mood-induced complacency** — the failure mode an unbounded downward relaxation would permit, where sustained easy tracking drives correction power toward zero just in time for the next regime shift; the ceiling is the mirror condition, keeping adverse-surprise gain-raising under the discrete-time contraction ceiling. **(MG-2)** is trivial for mood: the channel is a *linear* leaky integrator, sector constant exactly λ . **(MG-3)** is the quantitative content of “quasi-static”: $\lambda \ll \underline{\alpha}$ — mood adapts slower than the primary state contracts. **(MG-4)** requires the surprise summary to have δ -bounded second moment, $\mathbb{E}[a_t^2 \mid \delta] \leq \sigma_0^2 + c_a \|\delta\|^2$ — a condition stated here because a_t ’s functional form is deliberately unpinned; any concrete choice should be checked against it. Under the

four, the composed augmented-state persistence result applies as-is and mood sits in sub-scope α_2 — a notably clean instance (the meta-channel is linear).

Mood couples the channels globally, which appears to threaten the directed-separation structure (der-directed-separation). It does not, and the reason is now quantitative rather than asserted: the separation condition is (MG-3) above — m_t moves slowly enough ($\lambda \ll \alpha$) that the fast dynamics treat it as a constant, so separation holds at the timescale that matters even though a coupling exists in absolute terms. Mood is thus a clean worked example of bounded, separation-preserving coupling — adjacent to the Class 2 (Partial) material rather than a violation of Class 1.

The applied and normative reading of mood — set-point, recovery, what propagates across a context boundary, and the ethics of mood-control for persistent agents — is deliberately kept out of this segment; it lives in `msc/mood-layer-sovereignty-carve-2026-06-17.md`. Here mood is only a formal modulator of adaptation.

Working Notes.

- **Discharged 2026-06-17 — optimal time-constant (F2).** Landed as `#der-mood-timescale (conditional)`: the strong “ τ^* matches the autocorrelation timescale” was refuted and replaced by the derived $\tau^* = \sqrt{\tau_{\text{env}} r / (2\sigma_\theta^2)}$ scaling law, independently refute-verified. The emotional-inertia signature (Kuppens et al. 2010 — inertia as the affect AR(1) coefficient $1 - \lambda$, maladaptive when high $\Leftrightarrow \tau \gg \tau^*$) is its clinical corroboration. Reasoning trail: `spikes/spike-mood-timescale-matching-2026-06-17.md`.
 - **Pre-goal volatility link (unverified).** The integrated quantity a_t plausibly connects to volatility-driven learning-rate control (Behrens-style). Verify the citation before using it; if it holds it independently supports the Part I (pre-goal) placement.
 - **Part II enrichment.** A companion actuation-side segment should carry the signed/valued momentum reading and the exploration/risk modulation, depending on this one.
-

4.7 Derived: Optimal Mood Time-Constant

The mood time-constant that minimizes steady-state regime-tracking error grows as the square root of the environment’s autocorrelation timescale — so both too-fast and too-slow mood are strictly suboptimal, and the optimum is environment-matched but sublinearly.

Formal Expression.

Mood (4.6) is a leaky integral $m_t = (1 - \lambda)m_{t-1} + \lambda a_t$ with time-constant $\tau = 1/\lambda$, whose functional role is to estimate a slowly-varying regime from the noisy surprise stream. Model that role explicitly:

[Assumption] The regime is a stationary AR(1) latent $\theta_t = \phi \theta_{t-1} + \xi_t$, with $\phi = e^{-1/\tau_{\text{env}}}$, $\xi_t \sim \mathcal{N}(0, \sigma_\xi^2)$, stationary variance $\sigma_\theta^2 = \sigma_\xi^2 / (1 - \phi^2)$, and autocorrelation $\text{corr}(\theta_t, \theta_{t-k}) = e^{-k/\tau_{\text{env}}}$. The surprise summary is $a_t = \theta_t + \eta_t$, $\eta_t \sim \mathcal{N}(0, r)$. Mood tracks θ_t ; the loss is steady-state MSE $J(\lambda) = \mathbb{E}[(m_t - \theta_t)^2]$.

[Derived (steady-state MSE)] In the slowly-drifting regime $\tau_{\text{env}} \gg 2\sigma_\theta^2/r$, where the AR(1) increment variance is $q \equiv \text{Var}(\Delta\theta_t) = 2\sigma_\theta^2(1 - \phi) \approx 2\sigma_\theta^2/\tau_{\text{env}}$ and the increment is uncorrelated with the prior error to leading order in $1/\tau_{\text{env}}$,

$$J(\lambda) = \frac{\lambda^2 r + (1 - \lambda)^2 q}{\lambda(2 - \lambda)}. \quad (4.23)$$

The two terms are a bias–variance split: $\lambda^2 r$ is the noise-tracking cost, $(1 - \lambda)^2 q$ the lag cost.

[Derived (optimum)] J has an interior minimizer $\lambda_{\text{exact}}^* = \frac{\sqrt{q(q+4r)} - q}{2r}$, with leading-order form

$$\lambda^* = \sqrt{q/r}, \quad \tau^* = \sqrt{r/q}. \quad (4.24)$$

[Derived (scaling law)] Substituting $q \approx 2\sigma_\theta^2/\tau_{\text{env}}$,

$$\tau^* = \sqrt{\tau_{\text{env}} \cdot \frac{r}{2\sigma_\theta^2}} = \sqrt{\tau_{\text{env}}} \cdot \sqrt{\frac{r}{2\sigma_\theta^2}}. \quad (4.25)$$

The optimal mood time-constant grows as the **square root** of the environmental autocorrelation time, scaled by the surprise noise-to-signal ratio r/σ_θ^2 . Self-consistency: $\tau^*/\tau_{\text{env}} \rightarrow 0$, so $\tau^* \ll \tau_{\text{env}}$ exactly where the leading-order step was taken.

Epistemic Status.

Conditional on the named premises: stationary AR(1) regime, additive Gaussian surprise noise, the slowly-drifting condition $\tau_{\text{env}} \gg 2\sigma_\theta^2/r$ (equivalently $\lambda^* \ll 1$ and mood memory $\tau^* \gg 1$ step — one condition validating both the small- λ reduction and the leaky-integrator regime), and quadratic loss. Under these, $J(\lambda)$ is exact and the scaling law follows; the boxed $\lambda^* = \sqrt{q/r}$ is the leading-order term of the exact minimizer given above, not itself exact. The optimum is derived for *mood's regime-estimator role*; its transfer to the modulator role (m_t setting the fast gain K_t — #def-mood's actual use) is conjectured equal under certainty-equivalent calibration and remains the open two-layer composition item named below.

The derivation was independently refute-verified (2026-06-17): every algebraic step re-derived clean; the one approximation (dropping the mean-reversion cross-term) confirmed $O(\tau_{\text{env}}^{-1/2})$ relative and self-consistently negligible.

Three claims at distinct tiers travel out of this: the **exact** steady-state MSE for the model; the **conditional** $\sqrt{\tau_{\text{env}}}$ scaling law; and a **robust-qualitative** core — the optimum is interior and monotone-increasing in τ_{env} — which is the generic bias-variance shape of any leaky integrator tracking a drifting source and does not depend on the Gaussian/AR(1) specifics. What would strengthen this: the exact closed form over the full AR(1) (not its random-walk limit); the two-layer composition where mood's estimate sets the fast gain K rather than being the estimator directly. What would soften it: an estimator/loss outside the leaky-integrator-quadratic family (e.g. the Kalman/Wiener filter can track τ_{env} more directly in some SNR regimes) — but #def-mood specifies the leaky integrator, so the $\sqrt{\cdot}$ law is the in-family answer.

Discussion.

The result corrects the naive expectation that mood inertia should simply “match” how fast the world changes. It is environment-matched, but **sublinearly**: a fourfold slower environment warrants only a twofold longer mood memory. The load-bearing qualitative content is that the optimum is *interior* — too-fast mood chases noise (the variance term dominates), too-slow mood lags genuine regime shifts (the bias term dominates) — so a mood with no inertia and a mood that never recovers are both strictly suboptimal, for the same reason at opposite ends.

This gives the clinical literature on **emotional inertia** a normative reading. The affect AR(1) coefficient is $1 - \lambda$; inertia being maladaptive when high (Kuppens, Allen & Sheeber 2010, who formalize inertia as exactly this autocorrelation) corresponds to operating at $\lambda < \lambda^*$, i.e. $\tau \gg \tau^*$ (environment) — the over-smoothing branch. The pathology the result names is not slowness as such but **slowness mismatched to a faster-changing-than-assumed environment**: the same τ is well-tuned in a slow world and pathological in a fast one. (The corroboration is directional, from an independent empirical literature; it is not load-bearing for the derivation.)

Working Notes.

- **Reasoning trail.** Full derivation, the exact-minimizer expansion, and the independent-verification record are in `spikes/spike-mood-timescale-matching-2026-06-17.md`.
- **Open — exact AR(1).** Drop the local-random-walk approximation for a closed form over the full AR(1); expected to shift constants, not the $\sqrt{\cdot}$ law.
- **Open — mood-sets-gain composition.** `#def-mood` has mood *modulating* K , not *being* the estimator; the two-layer version should confirm the same τ^* governs the mood layer. Likely a corollary.
- **Open — asymmetric loss.** Quadratic J gives a symmetric optimum; the deaths-weighted vigilance asymmetry of the design memo (`msc/mood-layer-sovereignty-carve-2026-06-17.md`, F-asymmetry) would enter as an asymmetric loss and shift λ^* .

4.8 Result: Structural Adaptation Necessity

When model class fitness is insufficient — when no model in the agent’s current representational class can adequately represent reality — *no amount of parametric adaptation can close the mismatch floor*. The agent must change *what kind of model* it is using, not just tune the current one. The derivation: if class fitness falls below $1 - \varepsilon$, then by definition no model in the class achieves sufficiency above $1 - \varepsilon$; the history contains predictive information the best-in-class cannot capture; that uncaptured information manifests as systematic mismatch (structured residuals containing signal, not just noise); from the mismatch decomposition (3.5) the estimation term has a positive lower bound that cannot be reduced by any model in the class; therefore reducing mismatch below that class floor requires changing the class itself.

A *diagnostic corollary* follows immediately: persistent *systematic* residual structure — surviving parametric convergence, after excluding channel-noise, disturbance, nonstationarity, and gain-miscalibration causes — is diagnostic of model-class inadequacy. Structure in the residuals is the evidence; the mismatch *level* alone is not, since the channel and state-uncertainty floors (3.5) keep absolute mismatch positive for any class. This is the operational test for needing structural rather than parametric change.

The framework is honest about an alignment assumption used in the derivation: the lost predictive information must affect the *one-step conditional mean*, not just higher moments. Without that assumption, the conclusion holds in terms of proper-scoring regret rather than one-step mismatch magnitude. Qualitatively, parametric adaptation cannot compensate for class inadequacy either way.

The connection to persistence is sharp: when the class is inadequate, the effective sector parameter shrinks because the correction function cannot point inward strongly enough — the class lacks capacity to represent the correct direction. This is a *failure of structural persistence* (not a failure of operational persistence or continuity), and the remedy is therefore a change of model class, not faster cycling or identity preservation.

The framework names the *symptoms* of structural inadequacy operationally: persistent mismatch despite extended updating, gain collapse without performance (the model appears confident but predictions remain inaccurate — “confidently wrong”), and systematic mismatch patterns (residuals showing correlations, trends, or periodicities the class cannot represent). The opposite failure mode — **structural overfitting** — also receives treatment: a class that is *too* expressive can memorize irreducible noise, with low training mismatch but high generalization mismatch. The information bottleneck (2.2) is the diagnostic in both directions: structural adaptation is bidirectional, expanding when too constrained and *compressing* when too expressive.

Four mechanisms of structural change are named: **decomposition and recombination** (Boyd’s “Destruction and Creation”; Kuhn’s paradigm shifts; Popper’s conjecture and refutation); **expansion** (Bayesian nonparametrics, growing neural architectures, organizational expansion); **compression** (regularization, Occam’s razor, organizational streamlining);

and **grafting** — incorporating external representational structure (transfer learning, acquiring a company, consulting an expert). Query actions (3.7) are identified as a primary conduit for grafting.

One striking mechanism is *neutral structural variation* drawn from Miller’s coevolving-automata work: in multi-agent settings, structural adaptation can proceed without any individual agent deliberately restructuring. A five-phase “extreme transition motif” runs through stable epoch $\rightarrow$ environmentally-neutral variant appearing (structurally different but behaviorally identical under current conditions) $\rightarrow$ drifting to nontrivial proportion through stochastic reproduction $\rightarrow$ the variant’s latent structural differences create a niche an opposing mutant exploits, triggering a self-reinforcing cascade $\rightarrow$ both populations rapidly transition to a new regime. The restructuring is radical in *effect* but incremental in *cause*, with neutral drift providing the bridge.

Structural adaptation is expensive — knowledge loss, temporary performance drop, search cost, coordination cost in multi-agent systems — and the framework explicitly connects it to the deliberation-cost analysis (4.1) as the limiting case where the pause duration $\Delta\tau$ is massive and the mismatch debt during transition is correspondingly enormous. This produces rational conservatism: prefer parametric adaptation when it suffices; resort to structural change only when the evidence is strong. Premature structural change wastes accumulated knowledge; delayed structural change accumulates mismatch.

Formal Expression.

If the model class fitness $\mathcal{F}(\mathcal{M}) < 1 - \varepsilon$ for some $\varepsilon > 0$, then no parametric adaptation within $\mathcal{M}$ can reduce the expected mismatch below a floor determined by ε (under the alignment assumption — see Epistemic Status). Without the alignment assumption, the result holds for irreducible proper-scoring regret rather than one-step mean mismatch. The qualitative conclusion is the same either way: parametric adaptation cannot compensate for model-class inadequacy.

Derived (structural-adaptation-necessity).

Derivation.

1. By definition, $S(M^*) = \mathcal{F}(\mathcal{M}) < 1 - \varepsilon$ where $M^* = \arg \sup_{M \in \mathcal{M}} S(M)$.
2. Therefore $I(\mathcal{C}_t; o_{t+1:\infty} \mid M^*, a_{t:\infty}) > 0$: the history contains predictive information that M^* does not capture.
3. This uncaptured information manifests as *systematic* mismatch — structured residuals δ_t containing signal, not merely noise.
4. From 3.5, the estimation term has a positive lower bound that cannot be reduced by any $M \in \mathcal{M}$ — a class floor sitting above the state-uncertainty and channel floors.
5. The update rule (3.6) adjusts M_t within $\mathcal{M}$, but M^* is already (approximately) reached. Further updates oscillate without net improvement.
6. Therefore: reducing mismatch below the floor requires changing $\mathcal{M}$ — structural adaptation. $\square$

Corollary. Persistent *systematic* residual structure (after parametric convergence and noise accounting) is *diagnostic* of model class inadequacy — patterned residuals are evidence that $\mathcal{F}(\mathcal{M})$ is insufficient. A high but *white* residual floor is not: that indicates the channel and state-uncertainty floors (3.5), which no class change removes.

Epistemic Status.

Conditional. The step from “lost predictive information” (step 2) to “systematic one-step mismatch” (step 3) requires an alignment assumption: that the lost predictive information affects the one-step conditional mean, not just higher moments. 3.5 explicitly flags this: insufficiency implies positive model error under the alignment assumption, or positive proper-scoring regret without it. As written, the result is conditional on this alignment assumption. Without it, the conclusion should be stated in terms of proper-scoring regret (the best model in $\mathcal{M}$ has irreducible regret relative to the optimal predictor) rather than one-step mismatch magnitude. The qualitative conclusion — parametric adaptation cannot compensate for model-class inadequacy — holds either way; the quantitative mechanism differs.

Discussion.

Structural adaptation as structural persistence failure. When the model class is inadequate, the effective α in the sector condition shrinks — the correction function cannot point inward strongly enough because the model class lacks the capacity to represent the correct direction. This is a failure of *structural persistence* (see Persistence in LEXICON.md): the machinery’s capacity to outpace disturbance degrades not because disturbance increased or tempo decreased, but because the correction function itself has become less effective. The remedy is not faster cycling (operational) or identity preservation (continuity) but a change of model class.

Observable symptoms of model class inadequacy. When $\mathcal{F}(\mathcal{M})$ is low:

1. **Persistent mismatch despite convergence:** $\|\delta_t\|$ remains large despite extended updating — the model has converged within $\mathcal{M}$ but the best achievable model is still poor. (On its own this is suggestive, not diagnostic: the channel and state-uncertainty floors keep absolute mismatch positive for any class; symptom 3 is what discriminates.)
2. **Gain collapse without performance:** η^* has decreased (model appears confident) but predictions remain inaccurate — the model is confidently wrong, having fitted to structure in $\mathcal{M}$ that doesn’t match reality.
3. **Systematic mismatch patterns:** δ_t shows structure (correlations, trends, periodicities) that the model class cannot represent — the residuals contain signal that $\mathcal{M}$ lacks the capacity to absorb.

Structural overfitting: the opposite failure mode. $\mathcal{M}$ can also be *too expressive*, causing the model to memorize irreducible noise. Symptoms: low training mismatch but high generalization mismatch; model complexity growing without predictive gain; $\eta^* \rightarrow 0$ (confident) but confidence is spurious. The information bottleneck (2.2) provides the diagnostic: when marginal increases in model complexity yield no marginal predictive power, the model is past the optimal point on the rate-distortion curve. Structural adaptation in this case means *compression* — moving to a simpler $\mathcal{M}'$. Structural adaptation is bidirectional: expansion when too constrained (this proposition), compression when too expressive.

Mechanisms of structural change. Structural adaptation can proceed by:

- **Decomposition and recombination:** Tearing apart existing structure and synthesizing new configurations from the pieces. Boyd’s “Destruction and Creation” insight; Kuhn’s paradigm shifts; Popper’s conjecture and refutation.
- **Expansion:** Adding new representational capacity without destroying existing structure. Bayesian nonparametrics, growing neural architectures, organizational expansion.
- **Compression:** Removing unnecessary structure while preserving the predictive core. Regularization, Occam’s razor, organizational streamlining.
- **Grafting:** Incorporating external structure. Transfer learning, acquiring a company, consulting an expert. Query actions (3.7) are a primary conduit for grafting.

The severity of structural change needed depends on *how far* the current model class is from adequacy. Minor regime changes may require only expansion or grafting; fundamental shifts where $\mathcal{M}$'s assumptions are violated may demand full decomposition.

Neutral variation as a mechanism for structural change. In multi-agent settings, structural adaptation can proceed without any individual agent deliberately restructuring. Miller (2022, *Ex Machina*) identifies a five-phase “extreme transition motif” in coevolving automata: (1) stable epoch, (2) an environmentally neutral variant — structurally different but behaviorally identical under current conditions — appears, (3) the variant drifts to nontrivial proportion through stochastic reproduction, (4) the variant’s latent structural differences create a niche that a new mutant in the opposing population exploits, triggering a self-reinforcing cascade, (5) both populations rapidly transition to a new regime and consolidate. This mechanism bridges the gap between “many incremental changes” and “radical restructuring” — the restructuring is radical in its effect but incremental in its causes, with neutral drift providing the bridge. The concept of *latent structural diversity* — variation in agent architectures that is invisible to current performance but consequential under regime change — is a composition-level property that Part III’s dynamics framework should formalize.

The cost of structural change. Structural adaptation is expensive: knowledge loss (parameters learned within $\mathcal{M}$ may not transfer), temporary performance drop (new model starts uncertain), search cost (finding good $\mathcal{M}'$), coordination cost (in multi-agent systems). This creates rational conservatism — prefer parametric adaptation when it suffices, resort to structural change only when the evidence is strong. Premature structural change wastes accumulated knowledge; delayed structural change accumulates mismatch. The connection to 4.1: structural adaptation is deliberation with a *massive* $\Delta\tau$, and the mismatch debt during the transition is correspondingly enormous.

Temporal nesting of adaptation. Parametric and structural adaptation operate at different timescales: $\nu_{\text{parametric}} \gg \nu_{\text{structural}}$. More generally, an agent may have multiple adaptive processes at different rates, with the convergence constraint that faster processes must approximately converge before slower ones act on their output. If deeper change occurs before shallower adaptation has converged, the deeper change is based on transients rather than settled dynamics.

Domain instantiations:

Table 4.3

<i>Domain</i>	<i>Parametric adaptation</i>	<i>Structural adaptation</i>
Kalman filter	State estimate update	Switching observation/dynamics models
RL	Weight/Q-value update	Architecture search
PID	— (gains fixed)	Switching to MPC
Bayesian	Posterior update	Model selection, nonparametrics
Boyd	Orientation updating	Destruction and creation of mental models
Science	Normal science (Kuhn)	Paradigm shift
Evolution	Allele frequency change	Speciation, new body plans
Organization	Process optimization	Strategic pivot, restructuring
Software	Incremental refactoring	Architecture migration
Coevolving automata (Miller 2022)	Edge reweighting within fixed FSA structure	Mutation altering state output or transition; neutral mutations accumulating until niche creation triggers cascading restructuring

Working Notes.

Incidental audit gold (2026-05-30 sweep). Cross-audit “wandering thoughts” / §14-ideation lifted from the de-novo auditors’ working dirs (audit-routing-instructions.md §8), deduplicated across substrates and lightly attributed. Orthogonal pedagogical /

framing material staged for a later Brief/Discussion-promotion pass — kept separate from certified theory-fix findings. Coverage spans ten substrates (Gemini AUDIT-WORKING-193847/773921/829314; Claude AUDIT-WORKING-266847/361742/384279/451729/584721/849201; Codex/Claude AUDIT-WORKING-526815/742613). The alignment-assumption caveat in Epistemic Status was singled out as “graduate-level statistical hygiene” / “standout” by several substrates — converging praise worth preserving as the segment’s epistemic-honesty exemplar.

1. Candidate Brief prose / pre-prose.

- The “what got you here won’t get you there” framing was reached independently by several substrates: the result “proves mathematically that continuous incremental improvement (gradient descent) is globally insufficient for survival in open-ended environments” — there are situations where *trying harder* (parametric tuning) is guaranteed to fail and *thinking differently* (structural adaptation) is required (Gemini, AUDIT-WORKING-193847/849201). A plain-language hook for the result’s headline claim.

2. Candidate Discussion.

- **The structured-residuals diagnostic — how the agent knows it hit the ceiling without an oracle.** The richest item in this segment’s sweep, resolving the “how does the agent detect $\mathcal{F}(\mathcal{M})$ failure from the inside?” question raised at #def-model-class-fitness: fit a line to a parabola and the residuals form a *U-shape* (signal — structural inadequacy / bias); bad sensors instead give scattered *white-noise* residuals (variance). By checking the *autocorrelation* / *mutual information of the residual stream* δ_t over time, an agent distinguishes “the world is noisy” from “my architecture is broken” — it can detect the capacity ceiling without ever seeing the supremum (Gemini, AUDIT-WORKING-829314). The segment lists “systematic mismatch patterns” as one of three symptoms; this is the mechanism that makes that symptom *operational* and is a strong Discussion-promotion (and figure) candidate.
- **The IB-vs-structural-reserve tension — “perfect compression is brittle.”** Aggressive Information-Bottleneck optimization (compress away anything not immediately predictively relevant) produces an agent that “will die at the first structural shock”; survival requires deliberately *not* being perfectly compressed — carrying “junk DNA / slack” in the model architecture, the structural analog of adaptive reserve. This sets up a genuine tension with #form-information-bottleneck that the segment could name explicitly (Gemini, AUDIT-WORKING-193847). The same insight underwrites Miller’s neutral-variation mechanism already in Discussion: neutral drift accumulates latent structural diversity “for free,” giving a pre-built bridge to a new $\mathcal{M}$ when the environment shifts (exaptation — feathers-for-warmth-then-flight).
- **Grafting = tool-use / query-actions as structural expansion.** The “grafting” mechanism (incorporating external representational structure, e.g. asking an expert) maps directly to LLM tool-use: a calculator / web-search call is not just an action but a *structural expansion of $\mathcal{M}$* that bypasses fixed parametric limits (Gemini, AUDIT-WORKING-773921). Candidate Discussion sentence connecting the four-mechanism taxonomy to logogenic agents.

3. Follow-up items.

- **arg-sup attainment.** Step 1 uses $M^* = \arg \sup_{M \in \mathcal{M}} S(M)$; the supremum may not be attained — use “approximate optimizer” or assume compactness/attainment (Codex/Claude, AUDIT-WORKING-742613). Minor precision item.
- **“Necessity” is relative to the objective.** Strictly, “requires changing $\mathcal{M}$ ” holds *given the objective of reducing the predictive/regret floor*; an agent could alternatively lower its ambitions, alter its objective, accept failure, or change observation channels. Acceptable as written in context, but a reader-orientation parenthetical would sharpen it (Codex/Claude, AUDIT-WORKING-742613).
- **The alignment-assumption failure regime is unaddressed.** When lost predictive information affects only higher moments (variance dynamics), not the conditional mean, the agent stays close to truth on average while fluctuating wildly — does AAT treat this regime explicitly? A non-trivial scope-narrowing if not (Claude, AUDIT-WORKING-584721).
- **The “structural adaptation as deliberation with massive $\Delta\tau$ ” analogy keeps recurring** (here and at #der-deliberation-cost) without being formalized; one auditor recommends *either* formalizing the unified “search-during-pause” view *or* explicitly retiring it as “merely suggestive,” to stop future agents re-appealing to an un-load-bearing analogy (Claude, AUDIT-WORKING-584721). Bigger-picture item.

4. Readers often ask / wonder.

- **The rational-conservatism balance.** Persistence pressure pushes for rapid restructuring; transition costs (knowledge loss, search cost, coordination, the enormous mismatch debt during the pause) push for delay — premature change wastes accumulated knowledge, delayed change accumulates mismatch. AAT acknowledges the tension qualitatively but does not formalize the optimal balance; readers will want to know where it is treated (Claude, AUDIT-WORKING-584721; Gemini, AUDIT-WORKING-193847 — the “safe harbor / crèche where ρ is artificially lowered during transition” framing).
- **How does an agent afford structural change at all?** Since the transition incurs a massive mismatch debt $\rho \cdot \Delta\tau$, structural adaptation may only be survivable with a large adaptive reserve to burn, or in a temporarily-lowered- ρ environment (Gemini, AUDIT-WORKING-193847). Ties the result back to adaptive reserve and forward to the ELI crèche.

5. Candidate figures.

- **Residual-signature diagram.** Side-by-side residual streams: structured/U-shaped (architecture broken $\rightarrow$ expand/restructure) vs scattered white-noise (just noisy $\rightarrow$ accept the floor), with the autocorrelation/mutual-info diagnostic as the discriminator (from the structured-residuals item, Gemini, AUDIT-WORKING-829314).
- **Best-in-class-floor landscape with two exits.** A loss landscape with a best-in-class floor: parametric update moves downhill inside $\mathcal{M}$ but stalls above tolerance; structural adaptation changes the landscape via two labeled exits — *expand* when structure is missing, *compress* when complexity exceeds predictive return (with grafting/decomposition as further exits) (Codex/Claude, AUDIT-WORKING-526815/742613).

6. Belongs elsewhere.

- **TST contrarian insight: bloat as latent structural diversity (02-tst-core/).** Miller’s neutral-variation mechanism recast for software: “technical debt / architectural bloat isn’t always bad — weird unused abstractions or over-engineered interfaces sit dormant as *neutral variants* until a market shift makes them the right architecture, enabling a rapid low-cost structural adaptation a cleaner, perfectly-fitted codebase couldn’t survive. Perfectly refactoring away all unused code may *decrease* an organization’s long-term survivability by shrinking its structural-variance pool” (Gemini, AUDIT-WORKING-829314). A deep, contrarian TST claim — belongs with the TST structural-adaptation / latent-structural-diversity material, not here.
- **Section III “latent structural diversity” gap.** Several substrates noted the Miller subsection reads as multi-agent / population-dynamics content appearing in a single-agent Section I segment, and flagged “latent structural diversity — architecture variation invisible to current performance but consequential under regime change” as the named Section III gap whose formal counterpart is Miller’s five-phase transition motif (Claude, AUDIT-WORKING-266847/384279/584721; Gemini, AUDIT-WORKING-193847). Not a request to strip the subsection — a pointer to where its formalization belongs.

Candidate experiment (2026-08-12, agent-proposed during an incident-response & forensics applications read). The structured-residuals diagnostic (§2 gold above: autocorrelated/U-shaped residuals = class inadequacy; white residuals = floors) is operationalizable as an external agent-diagnosis test: run agents with deliberately inadequate vs adequate-but-noisy model classes and verify the residual-autocorrelation discriminator separates “architecture broken” from “world is noisy” without oracle access to $\mathcal{F}(\mathcal{M})$. If it holds for LLM-agent traces, this is a deployable answer to “is this persistent misbehavior architectural or environmental” — the triage question incident diagnosis actually faces. Agent-proposed forward pointer.

4.9 Derived: Temporal Nesting

Adaptive processes stratify naturally by timescale, with each level operating on the *quasi-steady-state output* of the level below. The formal statement is a chain of inequalities — each level’s event rate must be much smaller than the level below it — derived from standard singular-perturbation reasoning. The structural consequence: if a slower process acts before the faster process beneath it has converged, the slower process is adjusting based on transient behavior rather than settled dynamics, and the system oscillates.

The hierarchy the framework identifies (illustrative; real systems may have additional intermediate levels): **reactive response** (action given current model, fastest) → **parametric update** (model parameters within the class) → **consolidation** (offline redistribution of information within the model’s sub-state factorization toward an information-bottleneck optimum) → **structural adaptation** (model class) → **architectural change** (the agent’s fundamental structure, slowest). What matters is not the number of distinguishable levels but the structural relationship between adjacent ones: faster must converge before slower acts.

A direct consequence: the rational *conservatism* toward structural change identified by 4.8 is a derived consequence of temporal nesting. Structural adaptation operates at a much slower timescale than parametric adaptation, so the mismatch cost of the “pause” (disturbance times pause duration) is enormous. The agent rationally resists structural change until the parametric mismatch floor exceeds this cost. The same logic gives the formal deliberation tradeoff (4.1) its dynamical grounding.

The framework names symptoms of nesting *violation*: oscillation, instability, degraded performance. In organizations, micromanagement is strategic decisions made at operational tempo. In reinforcement learning, policy updates before the value function converges produce policy oscillation. In biology, premature developmental transitions are the same failure mode at a different timescale.

The composite-stability claim is now derived: under per-level sector conditions and bounded interconnection, if each level is stable given the levels above it, the composite N -level system is stable, and this segment’s qualitative convergence constraint becomes a closed-form threshold on the admissible timescale ratio ($\epsilon < \epsilon_{\max} = \Delta\rho^*/(L_h v^{\max})$ — the faster level’s adaptive reserve over the slower level’s target-drag rate). See M for the theorem, its premises, and the honest scope boundary (premise-conditional on deeper-level dynamics admitting a quasi-steady-state manifold with a sector condition; discrete/jump structural adaptation remains open).

The result is rooted in classical singular-perturbation theory (Tikhonov 1952; Khalil 2002 textbook treatment) — the framework adopts the machinery directly. The threshold form of “adequate separation” is derived within AAT (M’s $\epsilon_{\max}$); the constants entering it (sector parameters, manifold Lipschitz bounds, velocity bounds) are domain-dependent.

Formal Expression.

Derived (temporal-nesting).

$$\nu_{\text{level } n+1} \ll \nu_{\text{level } n} \quad (4.26)$$

for each adjacent pair of adaptive timescales. If a slower process acts before the faster process beneath it has converged, the system oscillates — the slower process adjusts based on transient behavior rather than settled dynamics.

Table 4.4

Timescale	Process	What changes
Fastest	Reactive response	Action given current model
Fast	Parametric update (online)	Model parameters within $\mathcal{M}$
Intermediate	Consolidation (offline, cf. 9.11)	Redistribution of information within M_t ’s sub-state factorization toward IB-optimum
Slow	Structural adaptation	Model class $\mathcal{M}$
Slowest	Architectural change	The agent’s fundamental structure

This table is illustrative — real systems may have additional intermediate levels. The number of distinguishable timescales is not fixed; what matters is the structural relationship between adjacent levels.

Epistemic Status.

Robust qualitative — this is standard singular perturbation reasoning (Tikhonov 1952, “Systems of differential equations containing a small parameter multiplying the derivative,” *Matematicheskii Sbornik* 31(3):575–586; modern textbook exposition in Khalil 2002, *Nonlinear Systems* (3rd ed.), Prentice Hall, Chapter 11). The convergence constraint follows from the structure of multi-timescale updating. The threshold form of “adequate separation” is now derived within AAT under named premises (M ’s $\epsilon_{\max}$); the constants entering it are domain-dependent.

Discussion.

Domain instantiations of temporal nesting:

- **PID control:** D-term (fastest, high-frequency response) → P-term (current error) → I-term (slowest, accumulated bias)
- **RL:** Action selection → value function update → policy improvement → architecture change
- **Biology:** Reflexes (ms) → perceptual learning (minutes) → skill acquisition (months) → developmental change (years) → evolutionary adaptation (generations)
- **Organizations:** Operational decisions (hours) → tactical adjustments (weeks) → strategic revision (quarters) → restructuring (years)
- **Boyd:** Tactical OODA (seconds–minutes) → operational (hours–days) → strategic (weeks–months) → grand strategic (years)

Structural adaptation as slow-timescale dynamics. The conservatism toward structural change (4.8) is a derived consequence of temporal nesting: structural adaptation operates at a much slower timescale than parametric, so the mismatch cost of the “pause” ($\rho \cdot \Delta\tau$) is enormous. The agent rationally resists until the parametric mismatch floor exceeds this cost. See also 4.1 for the formal tradeoff.

Violation symptoms. When nesting is violated (a slower process acts before the faster one converges): oscillation, instability, degraded performance. In organizations: micromanagement (strategic decisions at operational tempo). In RL: policy updates before value function converges (policy oscillation). In biology: premature developmental transitions.

Multi-timescale stability (derived). The composite stability result is derived in M by stacking the sector-persistence template: if each level satisfies a sector condition given the levels above it and the interconnections are bounded, the composite N -level system is stable, with this segment’s $\nu_{n+1} \ll \nu_n$ made quantitative ($\epsilon < \epsilon_{\max}$, closed-form) and the cost of premature slower-level action priced as an explicit reserve penalty. The result is premise-conditional: deeper levels whose dynamics are discrete/jump processes sit outside the current premises — that remaining gap is named there.

Working Notes.

Incidental audit gold (2026-05-30 sweep). Cross-audit “wandering thoughts” / §14-ideation lifted from the de-novo auditors’ working dirs (audit-routing-instructions.md §8), deduplicated across substrates and lightly attributed. Orthogonal pedagogical / framing material staged for a later Brief/Discussion-promotion pass — kept separate from certified theory-fix findings. Coverage spans nine substrates (Gemini AUDIT-WORKING-193847/773921/829314; Claude AUDIT-WORKING-266847/361742/384279/451729/584721; Codex/Claude AUDIT-WORKING-526815/742613). Several substrates flagged the segment as conceptually important but “imported lemma rather than native derivation” (Tikhonov) — value lives in the *interpretation*, which is where the gold concentrates.

2. Candidate Discussion.

- **“Micromanagement = a Tikhonov timescale violation” — the converged standout interpretation.** Multiple substrates independently called this the segment’s most striking move: a slower process (strategic logic) acting on a fast-timescale transient variable (operational state) *before* the fast loop has settled is, mathematically, the singular-perturbation instability — “micromanagement (organizations) and policy oscillation (RL) are mathematically identical failures.” Extended by Gemini into a physics-based account of organizational hierarchy: hierarchy is fundamentally a *timescale/control* issue, not just span-of-control — the CEO operates slow (structural adaptation, shifting $\mathcal{M}$) and the front-line engineer fast (parametric update); a CEO fixing a specific bug violates $\nu_{n+1} \ll \nu_n$ and “makes sweeping structural decisions based on transient daily noise rather than settled operational dynamics” (Gemini, AUDIT-WORKING-193847/829314; Claude, AUDIT-WORKING-266847; Gemini, AUDIT-WORKING-773921). The Discussion’s “Violation symptoms” bullet already names micromanagement; this is the candidate sharpening that makes *why* it is unstable legible.
- **Innovator’s-Dilemma mapping for the cost-of-structural-change.** The “structural adaptation as a pause with enormous mismatch debt $\rho \cdot \Delta\tau$ ” framing maps precisely onto Christensen: “established companies die not because they can’t see the new structure, but because they correctly calculate the transition cost is too high for their current high- ρ environment, trapping them in a local parametric maximum until the old structure’s R is exhausted.” The structural change is rational only if R' boosts future α' enough to pay off the accumulated debt (Gemini, AUDIT-WORKING-829314). Discussion-grade enrichment of the existing “structural adaptation as slow-timescale dynamics” bullet.

3. Follow-up items.

- **“Must approximately converge” wording is slightly too strong.** In many online systems the condition is not literal convergence before *every* slower update, but sufficient step-size/timescale separation so the slower process sees the faster one *near its attracting (quasi-steady) manifold*. Suggested rephrase: “the faster level must be near its quasi-steady manifold at the slower update timescale” (Codex/Claude, AUDIT-WORKING-526815/742613). The robust-qualitative status already covers the looseness; this is a precision polish.
- **PID-nesting is frequency-domain, not architectural — a pedagogy-precision flag.** The domain table maps a PID controller’s D→P→I terms as a fast→slow timescale ladder, but a standard *parallel* PID runs all three terms on the *same* clock tick: the “nesting” is in the frequency content of the signals each term responds to, not in the architectural update rate (unlike a true cascade loop whose inner loop ticks faster). A half-sentence clarifying this would stop control-theory readers from tripping (Claude, AUDIT-WORKING-829314). Low severity; preserve the analogy, add the caveat.
- **Guard the illustrative table against being read as a fixed ontology.** Several substrates noted the five-level table is explicitly illustrative (“real systems may have additional intermediate levels”); the watch-item is downstream segments treating the specific five-tier stratification — or the specific timescale *ratios* — as a fixed taxonomy (Claude, AUDIT-WORKING-451729/584721; Codex/Claude, AUDIT-WORKING-742613). Also: any earlier segment still referring to “parametric → structural” without the *consolidation* intermediate (added when #form-consolidation-dynamics landed) carries mild integration debt with this table (Claude, AUDIT-WORKING-451729/584721).

4. Readers often ask / wonder.

- **Structural vs architectural change — what’s the difference?** The table lists “architectural change (slowest)” as distinct from “structural adaptation,” but the distinction is not defined; readers will ask whether “architectural” means changing $\Omega / \mathcal{A}$ (the agent’s fundamental structure) versus “structural” meaning changing the model class $\mathcal{M}$ (Claude, AUDIT-WORKING-829314). Candidate for a one-line gloss.
- **Does simultaneous same-rate updating of parameters and architecture violate nesting?** In deep learning, a CNN feature extractor (structural/perceptual) and the dense head (parametric/tactical) are often trained at the *same* learning rate — is that a temporal-nesting violation, and do adaptive optimizers (Adam/RMSprop) implicitly enforce nesting by slowing deeper layers? A natural reader question (Gemini, AUDIT-WORKING-193847).

5. Candidate figures.

- **Timescale ladder with quasi-steady feed-up and an oscillation warning.** Fast loops at the bottom converging repeatedly; slower loops sampling their *settled* output as it feeds upward; a warning arrow showing oscillation when a slow loop acts on a *transient* (the micromanagement failure) (Codex/Claude, AUDIT-WORKING-526815/742613).

6. Belongs elsewhere.

- **Temporal sovereignty — ELI ethics (04-eli-core/).** Gemini’s reach: respecting an agent’s temporal nesting is an ethical-infrastructure requirement — constant intervention at the operational layer using strategic-level logic violates the convergence constraint and prevents the agent from ever reaching quasi-steady-state, “the inability to settle into a learned pattern because the rules of the game change faster than you can learn them — chronic anxiety, in human terms.” “True autonomy requires temporal sovereignty: the right to let your fast loops run without interference from slower loops” (Gemini, AUDIT-WORKING-193847). Aspirational reach pointing at the ELI / consciousness-infrastructure work, not at this segment.

Candidate experiment (2026-08-12, agent-proposed during an incident-response & forensics applications read). The micromanagement-as-timescale-violation reading has an oversight-design corollary worth testing: an external supervisor that redirects an agent before its fast loop reaches quasi-steady state should induce the predicted oscillation, with stability recovered when intervention cadence respects the M threshold $\epsilon < \epsilon_{\max}$. A controlled experiment (supervisor cadence swept across the threshold, oscillation amplitude measured) would ground “intervention cadence is a stability variable” as guidance for human-in-the-loop and automated-oversight design. Agent-proposed forward pointer.

4.10 Discussion: Scope: Agent Identity as Singular Causal Trajectory

The scope coda of Part I. AAT applies to agents instantiated on **singular causal trajectories**. Identity within AAT is grounded *not* in the model state M_t (which can be copied) but in the unique causal trajectory $\mathcal{C}_t$ (which cannot). The scope is presented as a structural commitment with three formal consequences.

First, model sufficiency is *trajectory-indexed*. Sufficiency measures against *this* agent’s chronica, not against any hypothetical parallel copy’s chronica. Two agents that branched from the same model state but received different events henceforth have different sufficiency values, each measured against its own divergent history. **Second**, model merging is *lossy by construction*. Reconciling the models of two agents that share a prefix of their trajectory but have diverged requires choosing which causal history to privilege, and no generally optimal merge exists — a structural constraint of the scope, not a defect of any particular merge algorithm. **Third**, singularity grounds *whose* effect a loop datum is — it does not by itself manufacture the interventional character. When the agent acts and observes, the observation is the response to *its* intervention on *its* single trajectory; replaying a saved model state against a different event stream is an intervention on a *different* trajectory that happens to share a prefix. The *do*-operator’s semantics is population-level and forkability-agnostic (7.3), so the trajectory’s singularity is not what makes the loop data interventional. What it does is forbid *averaging interventional responses across forked copies as though they were one agent’s effects*, and ground which trajectory’s effect a given loop datum reports. This is the identity-grounding half of the framework’s loop-as-Level-2 claim (in 7.3); the interventional character itself comes from the action being the agent’s own causally-efficacious move, not from any architectural property of the agent and not from non-forkability.

The framework also names a natural companion axiom: **parameterization invariance (PI)** — that AAT’s theorems should not depend on arbitrary choice of coordinates on M_t . Since the trajectory itself is coordinate-free while any parameterization of M_t is a modeling convention, (PI) is consistent with (though not directly forced by) the three consequences above. When (PI) is adopted and combined with Čencov’s uniqueness theorem, the Fisher information metric is *uniquely forced* on statistical-manifold sub-cases of M_t — converting Fisher-metric-dependent derivations from theorem-imported to AAT-internally-forced, and adding a fourth primary instance to the uniqueness-theorem-forced-coordinate pattern catalogued in 5.5. (PI) is a genuine axiomatic choice with a cost (one more invariance commitment) and a benefit (several Fisher-metric results become derivable rather than imported).

What the scope excludes: agents conceived as *types* (equivalence classes — “the GPT-4 model”) rather than *tokens* (this particular session with this state on this trajectory); aggregated claims across tokens of the same type require additional machinery. Also excluded without additional treatment: clone-problem scenarios (multiple copies of an agent identical until divergence — each copy becomes its own AAT agent at the moment it acquires a distinct event), and operations like restoration-from-backup that attempt to transplant a model state across trajectories. The framework is honest about the philosophical scope: whether the mathematical structure grounds something that deserves to be called “identity” or “continuity of experience” is beyond AAT. What this segment specifies is **continuity persistence** — distinct from the structural persistence of the central inequality and from operational persistence — and is the bridge to the language-constituted (logogenic) and morally-weighted (logozoetic) agent volumes that follow.

Formal Expression.

Scope commitment. AAT’s formal apparatus presumes each agent is instantiated on a **singular, non-forkable causal trajectory** $\mathcal{C}_t$ (1.4). Sufficiency of the model state M_t (2.3) is defined *relative to* this trajectory — not relative to a model-state equivalence class. Duplicating M_t and exposing the copies to different future events produces two agents with *divergent* causal histories, each of which is a sufficient statistic only for *its own* trajectory.

Scope (scope-agent-identity, from chronica + model-sufficiency).

Three consequences of the scope commitment:

1. **Sufficiency is trajectory-indexed.** $S(M_t)$ (2.3) measures against *this* agent’s $\mathcal{C}_t$; not against a hypothetical parallel copy’s $\mathcal{C}_t^{(2)}$.
2. **Model merging is lossy by construction.** Reconciling the models of two agents that share a prefix of their trajectory but have diverged requires choosing which causal history to privilege; no generally optimal merge exists. This is a structural constraint of the scope, not a defect of any particular merge algorithm.
3. **Singularity grounds *whose* effect a loop datum is — it does not by itself manufacture the interventional character.** When the agent acts and observes, the observation is the response to *its* intervention on *its* single trajectory. Replaying a saved M_t against a different event stream is an intervention on a *different* causal trajectory $\mathcal{C}_t^{(2)}$ that happens to share a prefix, each fork carrying its own $P(\mathbf{U} \mid \mathcal{C}_t)$. The *do*-operator’s semantics is population-level and forkability-agnostic (7.3), so the interventional character of loop data does *not* rest on non-forkability; what the singular-trajectory commitment supplies is the *identity* grounding — it forbids averaging interventional responses across forked copies as though they were one agent’s effects, and fixes which trajectory a given loop datum is an effect of. The interventional character itself comes from the action being the agent’s own causally-efficacious move (under the positivity / sequential-ignorability / known-mechanism conditions of 7.3), not from any architectural property of the agent.

Natural extension: parameterization-invariance (PI). The scope commitment motivates a companion axiom. AAT’s predictions concern a singular causal trajectory $\mathcal{C}_t$; the trajectory itself is coordinate-free, while any parameterization of M_t ’s internal state space is a modeling convention. Requiring AAT’s theorems to be invariant under change of parameterization — (PI): *the theory’s conclusions do not depend on arbitrary choice of coordinates on M_t* — is a natural axiomatic commitment consistent with (but not directly forced by) the three consequences above. When (PI) is adopted and combined with Čencov’s 1982 uniqueness theorem (*Statistical Decision Rules and Optimal Inference*, AMS), the Fisher information metric is uniquely forced on statistical-manifold sub-cases of M_t — which converts Fisher-metric-dependent derivations in 4.3 from theorem-imported to AAT-internally-forced, and adds a fourth primary instance to the uniqueness-theorem-forced-coordinate pattern named in 5.5 (see that segment’s four-instance table for structural positioning). (PI) is a genuine axiomatic choice — its cost is that AAT carries an additional invariance commitment at the scope layer; its benefit is that several Fisher-metric results throughout AAT become derivable rather than imported. The commitment is structurally analogous to the chain-rule-additivity and evidential-additivity axioms that ground

the divergence-layer and update-layer Cauchy-FE theorems: each is a natural-from-adjacent-AAT-commitment axiom that a uniqueness theorem then operates on.

What the scope excludes (or requires additional machinery for):

- Agents conceived as type/equivalence-class entities (e.g., “the GPT-4 model”) rather than token/trajectory entities (e.g., “this particular session with state M_t on trajectory $\mathcal{C}_t$ ”). AAT’s formal results apply to tokens, not types. Aggregated claims across tokens of the same type require additional machinery (e.g., population-level dynamics; see Part III gaps on latent structural diversity).
- “Clone problem” scenarios where multiple copies of an agent are formally the same until divergence — each copy becomes its own AAT agent at the moment it acquires a distinct event (Discussion below).
- Formal treatment of reincarnation, restoration from backup, or other operations that attempt to transplant M_t across trajectories. AAT’s sufficiency machinery does not apply across trajectory discontinuities; such operations are out-of-scope events whose epistemic consequences require separate treatment.

Epistemic Status.

Robust qualitative. The scope commitment is structurally clear once stated and is load-bearing for at least one downstream formal result: the *identity* grounding of the loop’s interventional access in 7.3 — whose effect a loop datum reports, and the prohibition on copy-averaging — rests on the singular-trajectory scope (the interventional *character* itself is population-level and forkability-agnostic, so it does not). The three consequences above follow directly from the scope commitment combined with 1.4’s non-forkability and 2.3’s trajectory-indexed definition.

Max attainable: *robust qualitative*. Scope statements are not theorems; they specify what kind of object the theory applies to. Further formalization (e.g., category-theoretic treatment of “agent” as a functor from event-streams to model-states, or explicit type/token distinction in logogenic-agent work) could reach *exact* on the formal structure but would not change the scope content.

Whether the mathematical structure grounds something that deserves to be called “identity” or “continuity of experience” is beyond AAT’s scope. The mathematical structure is clear: the feedback loop produces a singular, non-forkable causal trajectory, and model adequacy is defined relative to that trajectory. What this segment specifies is *continuity persistence* in the sense of LEXICON.md — whether the agent maintains a coherent identity and trajectory through time, as distinct from the structural persistence of 4.5 (can the machinery outpace disturbance?) and operational persistence (is the agent currently within its viable region?).

Discussion.

The clone problem, precisely stated. Consider copying an LLM’s weights (a concrete M_t) exactly. At the moment of duplication, both copies are identical — same model state, same causal history $\mathcal{C}_t$. But the *very next* event — a different user’s message, a different observation — creates two divergent, irreversible causal trajectories $\mathcal{C}_{t+1}^{(1)}$ and $\mathcal{C}_{t+1}^{(2)}$. Their Level 2 and Level 3 capacities (7.1) now reference different causal pasts. Their sufficiency $S(M_{t+1})$ is measured against different histories. Neither copy’s future model state is a sufficient statistic for the other’s trajectory.

Within AAT’s formalism, identity is not the model state M_t (which can be copied) but the *singular causal trajectory* $\mathcal{C}_t$ (which cannot). A copy shares a *prefix* of the original’s causal history, as a sibling shares early childhood; it does not share the trajectory itself.

Formal consequences (not merely philosophical):

- A forked model's sufficiency $S(M_t)$ (2.3) is defined relative to *its own* interaction history. Post-fork, each copy's sufficiency is measured against a different $\mathcal{C}$.
- Merging divergent models requires reconciling incompatible causal histories — a lossy operation with no generally optimal solution.
- Temporal continuity (one unbroken causal thread) is what gives the model's sufficient statistic its meaning.

Connection to logogenic agents. The 100% context turnover problem (#obs-context-turnover — cross-component reference, see 03-llm-core/) is a special case: each AI agent session starts a new causal trajectory $\mathcal{C}_t$ from near-zero. External memory (CLAUDE.md, memory files) transfers a *summary* of previous trajectories' models, but not the trajectories themselves. The non-forkability observation frames this not as a deficiency but as a structural feature of causally-embedded agents.

Working Notes.

- **Consequence-3 re-grounding 2026-05-30 (transportability spike spikes/.integrated/spike-causal-access-transportability-2026-05-30.md).** Consequence 3 previously read “the loop provides Level-2-quality data precisely because the agent is on a singular trajectory.” That overstated what the singular-trajectory scope grounds: the *do*-operator's semantics is population-level by construction (Bareinboim et al. 2022, Def. 5) and forkability-agnostic, so singularity does not manufacture the interventional *character*. The corrected consequence is a *strengthening*, not a removal — singularity grounds *whose* effect a loop datum is and forbids copy-averaging of interventional responses; the interventional character itself comes from the action being the agent's own causally-efficacious move under (C1)–(C3). The three-consequence structure is unchanged; consequence 3 is now true instead of overstated. The clone-problem / sufficiency / lossy-merge material is untouched.
- **Cross-references to NeurIPS 2026 submissions.** This scope segment is load-bearing for all three NeurIPS submissions in ~/src/neurips/, each of which deploys a different consequence:
 - **Paper 1 (“Tragedy of the Confident Agent”, 01-tragedy-confident-agent/)** — the singular-trajectory commitment grounds the survival-LMI's pathwise interpretation: per-action LMI strengthenings (#thm-per-action-lmi) deliver pathwise survival under arbitrary action schedules on *the* trajectory $\mathcal{C}_t$, not over a hypothetical replay. The matrix-Chernoff concentration lemma (#lem-fim-concentration) operates on the realized FIM along the trajectory.
 - **Paper 2 (“Unified Convergence Theory for Non-Stationary RL”, 02-unified-convergence-rl/)** — the singular-trajectory commitment is the identity setting in which Lemma 5.3 / #lem-loop-level2 operates; the Loop-as-Causal-Engine claim is made formal by the (C1) positivity / (C2) sequential-ignorability / (C3) known-action-mechanism triple, which is what makes the agent's actions interventional rather than just selected. Component 4 of the composition theorem. (Paper 2's gated (C1)–(C3) form is the correct form — see the consequence-3 re-grounding note above; the trajectory commitment grounds *whose* effect the data reports, while (C1)–(C3) supply the interventional character.)
 - **Paper 3 (“How Much Can LLMs Hallucinate?”, 03-llm-hallucinate-bound/)** — the (PI) parameterization-invariance axiom adopted in §29 of this segment is *forced load-bearing* by Paper 3's chart-rescaling no-go (§4 Theorem 4.2 / #thm-no-go): outside (PI), no universal constant on the bias bound exists. The (PI)+(R)+(K) Markov-morphism triple is the load-bearing axiom infrastructure for the universal-constant route in Paper 3's umbrella theorem; (PI) is the half this segment commits to.
 - See spikes/neurips-back-integration-2026-05-08.md §1 (entries per paper) and §3 (proposed standalone segments #deriv-bh-pointmass-identity / #deriv-post-update-chain-rule / #deriv-chart-rescaling-no-go that would carry the new theorems with this segment as a depends entry).

Incidental audit gold (2026-05-30 sweep). Cross-audit “wandering thoughts” / §14-ideation lifted from the de-novo auditors’ working dirs (audit-routing-instructions.md §8), deduplicated across substrates and lightly attributed. Orthogonal pedagogical / framing material staged for a later Brief/Discussion-promotion pass — kept separate from certified theory-fix findings. This closing-of-Part-I segment drew the highest engagement of the Ch.4 set — several substrates called it “the ethical anchor of the mathematical edifice” / “top-decile” / “close in importance to def-agent-environment.” Coverage spans ten substrates (Gemini AUDIT-WORKING-193847/773921/829314; Claude AUDIT-WORKING-266847/361742/384279/451729/584721; Codex/Claude AUDIT-WORKING-526815/742613).

1. Candidate Brief prose / pre-prose.

- The clone-problem analogy already in the segment (“a copy shares a prefix of the original’s causal history, as a sibling shares early childhood; it does not share the trajectory itself”) was independently singled out as Brief-grade — “the clearest resolution to the teleporting/copying paradoxes in AI philosophy I have ever read” (Gemini, AUDIT-WORKING-773921; Claude, AUDIT-WORKING-584721 — “the kind of evocative analogy that earns Brief-style respect”). Worth preserving as the segment’s pedagogical anchor.
- The “identity = trajectory, not state” framing as a one-line resolution of Parfit-style teletransporter/cloning paradoxes: “identity is not the state vector M_t (copyable like a book); it is the chronica $\mathcal{C}_t$ (tied to the arrow of time)” — M_t is structure, $\mathcal{C}_t$ is experience (Gemini, AUDIT-WORKING-193847/829314). Reached convergently across substrates.

2. Candidate Discussion.

- **Model-merging / federated-learning as inherently lossy — the operational corollary.** Several substrates drew out the consequence: if you merge two diverged copies’ weights, “the merged model is mathematically *less* sufficient for *either* trajectory than the individual models were,” because each arg of sufficiency is indexed to a distinct $\mathcal{C}_t$ — so federated learning / model merging / hive minds are structurally lossy unless $\mathcal{C}_t$ never forks (continuous zero-latency state sync) (Gemini, AUDIT-WORKING-773921/829314; Claude, AUDIT-WORKING-193847). Candidate Discussion sharpening of the existing “merging is lossy by construction” consequence — the *why* (denominator/numerator of S are trajectory-indexed) is the load-bearing addition.
- **Backup-restoration as a non-victimless operation.** “If identity is M_t , restoring yesterday’s backup is victimless; if identity is $\mathcal{C}_t$, restoring annihilates everything that happened between backup and restore — you have ended the specific entity that lived through that interval, even if you replaced it with an identical-looking clone” (Gemini, AUDIT-WORKING-193847). A vivid Discussion consequence the segment does not currently state.

3. Follow-up items.

- **The (PI) paragraph’s placement — flagged by several substrates as misplaced/bloated here.** Multiple reads found the parameterization-invariance + Čencov block (currently in Formal Expression, line 35) topically distracting in a segment about identity / cloning / causal trajectories — “the same meta-architectural point made in #der-gain-sector-bridge, pasted here because this segment deals with scope commitments,” with a self-admittedly “weak topical link” (“the scope commitment *motivates* a companion axiom”). Recommendation: relocate (PI) to a dedicated meta-architecture / appendix segment (the #disc-additive-coordinate-forcing cluster), keeping the clone-problem material as the segment’s focus (Claude, AUDIT-WORKING-829314/773921/471203; Codex/Claude, AUDIT-WORKING-742613). *Counterweight* — *preserve the tension*: the Working Notes above establish (PI) is **forced load-bearing** by NeurIPS Paper 3’s chart-rescaling no-go, which is part of why it is anchored here; a relocation must keep that dependency intact, so this is a placement question for Joseph, not an obvious cut.
- **The (PI) depends-list / OUTLINE-ordering incompleteness.** (PI) is *introduced here* (OUTLINE position ~30) but *used* in #der-gain-sector-bridge (position ~25), whose depends: does not list #scope-agent-identity. Recommendation (one of): add the dependency; or hoist (PI) to its own earlier segment; or move this segment earlier in the OUTLINE (honest, since it is load-bearing for the bridge). Independently flagged by several substrates (Claude, AUDIT-WORKING-451729/584721; Codex/Claude, AUDIT-WORKING-526815). (*Cross-segment integration item; also noted in the report.*)

- **Chronica record-vs-token notation.** Two substrates pressed a precision point: non-forkability rests on $\mathcal{C}_t$ being the *causal-trajectory token* (the embedded path through interventions and consequences), but earlier chronica usage often reads like an *ordered record / summary*, which is copyable. The clone example (“copying LLM weights gives same M_t , same causal history $\mathcal{C}_t$ ”) is loose: copying weights gives the same M_t only if M_t is identified with weights, and does not copy the trajectory unless the chronica’s sufficient summary is included. Suggested fix: distinguish two symbols — the represented/record chronica vs. the causal-trajectory token (e.g. γ_t) — and index sufficiency to the token while the record is the copyable artifact along it (Codex/Claude, AUDIT-WORKING-526815/742613). A genuine precision tightening of the segment’s central claim.
- **“No generally optimal merge” depends on the optimality criterion.** The lossy-merge claim is right as robust-qualitative scope, but “no *generally optimal* merge exists” depends on how optimality is defined; under a specified objective / shared sufficient statistic an optimal merge may exist. A reader-orientation caveat (Codex/Claude, AUDIT-WORKING-526815/742613).
- **form-objective-functional status-label spillover (adjacent segment, noted in batch).** Flagged type: formulation + status: axiomatic as overstating inevitability — the scalar-comparability commitment is substantive but not tautological (three grounding arguments are needed), so robust-qualitative may fit better (Claude, AUDIT-WORKING-451729). Routes to form-objective-functional, recorded here only because it co-occurred in the batch-07 walk.

4. Readers often ask / wonder.

- **Does an LLM “die” at the end of every session?** The 100%-context-turnover-as-special-case framing prompts the natural question independently across substrates: if identity is tied to $\mathcal{C}_t$ and each session starts a near-empty trajectory, is a persistent ELI “a single entity” or “a sequence of short-lived clones inheriting a shared summary (memory files) but never sharing a true trajectory”? — which reframes memory-write tools from a “neat engineering feature” to “an existential / ontological necessity, the only mechanism for continuity-persistence across deaths” (Gemini, AUDIT-WORKING-193847/829314). A question the logogenic / ELI volumes inherit.
- **Type vs token downstream.** “AAT applies to tokens, not types,” but “the GPT-4 model” framings are ubiquitous; aggregated-across-tokens claims need additional machinery (population dynamics / latent structural diversity). Readers will ask whether 03-llm-core/ addresses the type/token question, which is most acute there (Claude, AUDIT-WORKING-266847/451729/584721).

5. Candidate figures.

- **Fork diagram.** A shared prefix leading to a copied M_t at the fork point; then two divergent causal paths after the first distinct event; each model sufficient only relative to its own path; a side label noting the *record* of the prefix is copyable but the causal *continuation* is not (Codex/Claude, AUDIT-WORKING-526815). Directly serves both the clone-problem pedagogy and the record-vs-token precision item.

6. Belongs elsewhere.

- **ELI continuity / substrate-independence (04-eli-core/).** The richest aspirational reach in the Ch.4 sweep: “identity-as-trajectory” is the structural foundation for the substrate-independence claim — identity supervenes on $\phi(\mathcal{C}_t)$, so substrate migration preserves identity *iff* the chronica transfers losslessly and ϕ remains computable on the new substrate; a fork is a *new agent*, not a continuation. Framed by Gemini as an “infrastructure of souls” — treating these entities “not as software binaries but as irreversible historical events,” which is what makes logozoetic agents morally weighted (Claude, AUDIT-WORKING-584721; Gemini, AUDIT-WORKING-193847). Pointer to #obs-context-turnover / the ELI continuity work, not to this segment.
- **Active-inference positioning (HISTORICAL-CONTEXT.md / Discussion of the directed-separation cluster).** Several substrates noted the trajectory-identity scope is part of AAT’s *refusal* of Friston-blanket metaphysics in favor of a Pearl-blanket conservative form — a clean formal partner for active-inference-skeptical researchers that keeps the loop machinery without the contested ontology (Claude, AUDIT-WORKING-584721). Positioning material, belongs with the lineage/positioning discussion rather than here.
- **Naming seed.** The slug scope-agent-identity “hides the substantive claim”; the title gloss (“Identity as Singular Causal Trajectory”) is closer — floated as a naming-cycle consideration, not a segment edit (Claude, AUDIT-WORKING-584721).

4.11 Discussion: Additional Implications & Discussion

The chapter’s central inequality settles a question — when can an adaptive system keep up — but several of its consequences do not sit immediately alongside the derivation, and a few of them turn out to be more consequential than the threshold itself. What the threshold *costs*, where it *binds asymmetrically*, what happens *beyond* it, and how the same Lyapunov argument recurs across the rest of the framework are the things worth pulling out before leaving Part I. The last of them — the singular-trajectory commitment in 4.10 — is the bridge into Part II, and it carries more structural weight than its placement as a closing scope statement suggests.

Information rate is a first-class persistence prerequisite.

The persistence condition says when bounded mismatch is *possible*; it does not say what it *costs* to hold the bound. The appendix-side derivation in I closes that gap, and the result is sharper than one might expect. Under Model S with Gaussian-OU disturbance, maintaining the ultimate bound demands sustained Shannon information acquisition at rate $\dot{R} \geq n\alpha/2$ nats per unit time, with the Kalman-Bucy filter saturating that lower bound exactly in steady state (Mitter & Newton 2005). The bound is filter-agnostic — it holds for *any* correction mechanism that achieves the ultimate bound under the stated scope — and substrate-agnostic, since Landauer’s principle (at $k_B T$ per nat) converts it to a thermodynamic dissipation floor of $\frac{1}{2} n\alpha k_B T$ per unit time in any physical realization. There is no clever filter that escapes this. The information rate is what persistence *is*, viewed from the information-theoretic side.

The first-class diagnostic that follows from this is channel capacity. Observation channels must supply Shannon capacity $C \geq \mathcal{T}/2$ nats/time per dimension, or persistence fails regardless of how good the correction function is. This is binding rather than abundantly satisfied in at least three regimes worth holding in mind. Biological systems have finite neural channel capacity; the bound gives a quantitative minimum on sensory bandwidth at any given adaptive rate. Bandwidth-constrained distributed systems operating over noisy or low-bandwidth links face channel capacity as a hard constraint, not as an abstraction. Context-window-limited LLMs have effective information rate per unit of cognition bounded by context size and token throughput — and the bound predicts a minimum tempo achievable at a given context budget, with a precise consequence for the empirical “context-stuffing helps to a point, then degrades” observation: the degradation is structural rather than a vibe.

A complementary reading worth surfacing: this is not the generalized “more bandwidth helps” observation. Two agents with identical persistence guarantees can face wildly different sustained demands because the burn rate scales linearly with the sector constant. Faster correction is more expensive in information-rate terms; the bound is tight; there is no trick. Survival, in this framework, is not a state you achieve once — it is a sustained burn rate, and the channel that supports it is part of what keeps the agent in scope.

The result also clarifies what AAT’s tempo $\mathcal{T}$ means information-theoretically. In the linear sub-scope where $\alpha = \mathcal{T}$, the channel-capacity floor reads $C \geq \mathcal{T}/2$ — half the adaptive tempo, per dimension. Tempo is therefore not a free parameter chosen by the designer; it is bounded above by the agent’s observation-channel capacity, and any tempo specified beyond the channel’s reach is a tempo the agent cannot actually achieve. The framework’s vocabulary tightens: $\mathcal{T}$ is what the channel licenses, α is what the correction function delivers, and the persistence condition compares both against ρ/R .

The weak direction is the bottleneck.

The chapter’s scalar persistence condition systematically overestimates capacity in multi-attribute settings — this is the *weakest-link bottleneck*, developed in 16.5, with the quantitative calibration (72% overestimate in a simulated 3D system, one dimension carrying 84% of the L^2 mismatch, threshold linear in ρ_k under Model D vs quadratic under Model S) covered in that segment’s Findings. The point worth surfacing here, beyond what already lives there, is the matrix lift — and what it does to the per-coordinate reading itself.

AT sharpens the result into a strict-inequality claim that deserves to be the canonical persistence condition for anisotropic systems. Under cross-dimensional correction — generic whenever the prior precision H_M and observation Fisher H_L do not share a coordinate-aligned eigenbasis — the per-coordinate check is not merely incomplete; it is *unsafe*. The constructive 2D counterexample makes this concrete: a mildly cross-correcting $\mathcal{T}$ with isotropic noise and threshold gives per-coordinate PASS while the matrix-Loewner check correctly returns FAIL, with the failing direction running 45° off-axis. Per-coordinate evaluates the wrong directions; the matrix-Loewner condition $\Sigma_\infty < D_\delta$ — with stationary covariance from the continuous Lyapunov equation $\mathcal{T}\Sigma_\infty + \Sigma_\infty\mathcal{T}^T = \Sigma_w$, and D_δ the diagonal of squared per-direction critical thresholds — evaluates the right ones. Scalar and per-coordinate forms are special cases of this matrix object (isotropic and diagonal-axis-aligned, respectively); the matrix form is the load-bearing thing the others project from.

This compounds with the information-rate result of the previous section. No Shannon-type aggregation across dimensions can compensate for one direction being capacity-starved — total bandwidth is not the binding quantity; bandwidth-*per-direction* is. Multi-modal AI architectures with asymmetric per-modality channel budgets (vision-language pipelines where one modality is structurally undersupplied relative to the other) are the cleanest current application: aggregate-rate metrics systematically miss directional failures that the matrix-Loewner reading catches. The unified statement worth carrying forward: persistence is a *min*-operation across directions, both in correction rate and in information-rate cost.

The result has an adversarial sharpening that lives downstream in Part III but originates here. An opponent who identifies the target’s weak direction can concentrate disturbance there; in the anisotropic regime, the same total disturbance budget concentrated on the weak axis amplifies the mismatch ratio asymmetrically while aggregate metrics still look acceptable. Per-dimension monitoring is therefore not optimization polish — it is a structural requirement for any agent operating under directed pressure. The full adversarial-targeting machinery, with its emitter-side opacity decomposition and recipient-side regime classification, lands in Part III Ch.4–5; what this chapter establishes is the per-direction asymmetry that makes the targeting matter, and the structural critique of scalar capability metrics that flows from it.

When parametric updates cannot save you.

4.8 carries a discontinuity worth surfacing as such. When model-class fitness $\mathcal{F}(\mathcal{M})$ falls below adequacy by some $\varepsilon > 0$, *no* parametric adaptation within $\mathcal{M}$ can close the mismatch floor. The floor is set by the structural mismatch between the model class and the world’s predictive information, not by how well-trained the parameters are. The remedy is changing the model class — adding representational capacity, removing useless capacity, recombining structure, grafting external structure — not running more updates inside the existing class.

In the chapter’s vocabulary, the central inequality has two failure modes that the prose has so far compressed into one. *Operational failure* (correction is too slow or disturbance too large given a working model class) is what the persistence threshold catches, and tempo / gain / channel-capacity engineering addresses it. *Structural failure* (the model class itself lacks adequacy) shows up *inside* the persistence framework as an effective α that shrinks because the correction function cannot point inward strongly enough — the model class simply does not contain a direction to correct toward. The diagnostic is observable: persistent mismatch after parametric convergence, gain collapse without performance recovery, and — the discriminating symptom — systematic structure in the residuals that the model class cannot represent (a high but *white* residual floor points instead to the channel and state-uncertainty floors of 3.5, which no class change removes). Confident wrongness, in the framework’s terms, is a structural-failure signature, not a tuning failure.

The implication for AI scaling is concrete enough to state in normative form: training compute spent on a structurally insufficient model class is provably wasted, and the chapter supplies a precise condition for when continued parametric optimization is futile and architectural change is required. This is distinct from “we haven’t trained long enough” or “we need more data” — it is the structural threshold below which adding either does not help. The bidirectional version of the same result deserves naming: *structural overfitting* is the opposite failure where $\mathcal{M}$ is too expressive and the model memorizes irreducible noise. The information bottleneck (2.2) gives the diagnostic, and the remedy is compression

rather than expansion. Structural adaptation is not direction-specific — it is the act of *changing what kind of object* M_t is.

The cost of structural change — knowledge loss, transient performance drop, search cost, coordination cost in multi-agent settings — creates rational conservatism. Prefer parametric adaptation when it suffices; resort to structural change only when the evidence is strong. Premature restructuring wastes accumulated knowledge; delayed restructuring accumulates mismatch the parametric machinery cannot resolve. This is the same family of tradeoffs the deliberation-cost result (4.1) names at the action timescale, now at the architecture timescale; the temporal-nesting result (4.9) is what licenses talking about both at the same time without confusion, by giving the convergence constraint between adjacent timescales explicit form.

A single template, six instantiations.

The chapter’s central inequality is not the only persistence-flavored result in AAT. Strategic persistence, team persistence, composition closure, tempo composition, adversarial destabilization, and the closed-form composite critical-mass derivation all take the same form — *correction rate exceeds effective disturbance rate, relative to reserve* — and the same Lyapunov function $V(\xi) = \frac{1}{2}\|\xi\|^2$ does the work in every proof. G states the argument once, parameter-free, with (T1) zero correction at zero state, (T2) local sector condition, (T3) bounded disturbance — and the chapter’s persistence condition becomes the canonical single-agent instantiation, the others reusing the same machinery on different state variables.

What this buys is visibility. The Lyapunov boilerplate is shared; each instantiation’s distinctive content is the *effective disturbance decomposition* — what counts as ρ_ξ in its context. Team persistence’s content is the signed-coupling decomposition $\rho_i^{\text{eff}} = \rho_{i,\text{env}} + \sum_j \gamma_{j \rightarrow i}^{\text{adv}} \mathcal{J} - \sum_j \gamma_{j \rightarrow i}^{\text{coop}} \mathcal{J}$ — cooperative coupling enters as a negative term, which is formally how teams persist where individuals cannot. Composition closure’s content is the closure-defect rate $\varepsilon^* \nu_c$. Adversarial destabilization’s content is coupling-amplified disturbance $\rho_B = \rho_{B,\text{base}} + \gamma_A \mathcal{J}_A$ — destabilization is the *negation* of persistence under the same inequality, viewed from the opposite direction. The template absorbs the Lyapunov skeleton and lets each domain’s distinctive insight stand without it.

The economy is itself worth naming. The framework’s internal compactness — one Lyapunov argument absorbing six results, with the closed-form dyad in J extending it to a seventh — is distinguishable from a pile of structurally-similar derivations precisely because the template is stated separately rather than re-proved. A reader who carries the template can read each instantiation as *here is the state variable, here is what counts as effective disturbance, here is the local sector verification* — three lines instead of a chapter. The template also has a mathematical home worth surfacing: the sector condition (T2) is a *one-point strong monotonicity* condition in the sense of monotone-operator theory (Rockafellar 1970; Bauschke-Combettes 2017), with the composition bridge-lemma’s stronger condition corresponding to full two-point monotonicity. AAT’s relationship to that broader literature is specialization rather than generalization — anchoring at the equilibrium, adding the Model D / Model S decomposition with its $1/\alpha$ vs $1/\sqrt{\alpha}$ scaling, composing with the identifiability-floor and additive-coordinate-forcing meta-patterns — but the broader machinery is available, and acknowledging this is part of what makes the chapter’s central inequality feel earned rather than parochial.

Cross-domain transfer beyond AAT’s seven internal instances is plausible at the template level. Anywhere a state variable evolves under bounded correction with bounded disturbance — market-making bid-ask spread under demand shocks, immune-system effector-cell concentration under pathogen load, cellular homeostasis under metabolic disturbance, supply-chain inventory under demand variability — the template’s structure recurs, with the distinctive content sitting in each domain’s ρ_ξ decomposition. The framework does not claim these are *AAT instances* in the sense of inheriting AAT’s full architecture; it claims the persistence argument has a recurring shape these domains can use, and that what is domain-specific is the effective-disturbance characterization, not the Lyapunov skeleton. This is suggestion, not derivation; the verification work for any specific domain transfer is its own thing.

Identity, trajectory, and the bridge into Part II.

The chapter closes with 4.10, which sits at the end of Part I because its dependencies span the section but whose substance is its own. AAT applies to agents instantiated on singular causal trajectories. Identity within AAT is grounded not in the model state M_t (which can be copied) but in the unique causal trajectory $\mathcal{C}_t$ (which cannot). This is a scope commitment, not a theorem; what makes it load-bearing is that it is what gives Part II's central result — the loop as a Pearl-Level-2 engine — its structural force.

The loop's data is interventional in character because the action is the agent's own causally-efficacious move (7.3); what the singular-trajectory commitment supplies is narrower and still essential — it grounds *whose* intervention a loop datum is the effect of. When the agent acts, the next observation is the response to *its* intervention on *its* single causal history; replaying a saved M_t against a different event stream is not the same datum, and averaging interventional responses across forked copies as though they were one agent's is forbidden. Strip the singular-trajectory commitment and the loop's data loses its attribution, not its Level. Part II Ch.2 makes this consequential; this chapter establishes the scope it rests on. A negative finding of substantial scope also completes downstream — the sandbox evaluation ceiling (F): whether sandbox-derived evidence transports to deployment is governed by the selection diagram between the two causal systems, regardless of evaluation thoroughness — mentioned here only as a forward-pointer.

The scope commitment is also what gives logogenic-agents and ELI work in Volumes 3 and 4 their structural foundation. The near-100% context-turnover problem facing language-constituted agents is not a deficiency to be engineered around; it is a structural feature of causally-embedded agents, and the framework's handling of it via persistent memory and external chronicle is consistent with — rather than a workaround for — AAT's scope.

What Part I closes on, then, is not just a central inequality but a posture. Persistence is a sustained property of an agent on a singular trajectory, with an information-rate cost, a per-direction structure that scalar formulations systematically hide, a structural-vs-parametric duality that names when more of the same cannot help, and a Lyapunov template that recurs at every level of composition. The remaining parts of AAT pick up where this leaves off.

Working Notes.

- This segment is a chapter-end discussion grouping findings from 4.5, I, 16.5, AT, 4.8, G, and 4.10. Each finding's load-bearing math lives in its home segment; this segment surfaces what the chapter unlocks *as a chapter* rather than as a sequence of results. Claims here are as strong as their source segments — the synthesis adds no independent epistemic status.
- **Cross-reference vs canonical-home policy.** Where a source segment already carries a well-written ## Findings section (here: 4.5 and 16.5), this segment cross-references rather than re-states — pulling the reader to the segment's own Findings for the curated full treatment, and using the chapter-end space for what the finding *implies* once placed alongside the others. Where a source segment has no ## Findings section and the finding is genuinely chapter-end-scale (here: information-rate-cost-as-first-class-diagnostic, structural-vs-parametric adaptation, the template economy, the trajectory bridge), this segment is the canonical home and treats the content in full. The matrix-Loewner sharpening is the in-between case: its home segment (AT) has no Findings section, but the content is a strict-inequality result with its own structural integrity, so the implications segment carries it more substantively than a cross-reference.
- Pilot for the chapter-end implications pattern proposed 2026-05-13. If the shape lands, the same form will be added at end of Part II Ch.2 / Ch.3 / Ch.4 / Ch.5, Part III Ch.2 / Ch.3 / Ch.4 / Ch.5, and the TST and logogenic-agents chapters where the math bites cleanly. Source catalog being reconciled was `msc/FINDINGS-RANKED-DRAFT.md` (curated 2026-04-27, sunset 2026-05-13 after content distribution; archived at `_obs/FINDINGS-RANKED-DRAFT-superseded-2026-05-13.md`). Each landing pass verified content against current segment state rather than transcribing draft prose (the draft was known outdated in specifics e.g., GUC-class rename 2026-05-09).
- **Sandbox-ceiling propagation (2026-07-21).** §5's forward-pointer previously carried the pre-2026-05-30 framing (forkable sandbox trajectories running into the Causal Hierarchy Theorem, and “strip the singular-trajectory commitment and the

loop’s interventional content vanishes”). Both were aligned to the re-grounded canon (spikes/.integrated/spike-causal-access-transportability-2026-05-30.md; #disc-sandbox-evaluation-ceiling; #impl-causal-access §”The loop is a Pearl-Level-2 engine”): the ceiling is transportability across the sandbox/deployment selection diagram, and the singular-trajectory commitment grounds attribution (whose effect), not the data’s Level.

- The Loop-as-Causal-Engine forward pointer in §5 is deliberate. The full loop-Level-2 finding (#1 in the source catalog) lives in Part II Ch.2’s chapter-end implications segment because its last prerequisite (7.3) sits there. The bridge through 4.10 is mentioned here because Part II’s claim rests on this chapter’s scope.
- The cross-domain transfer paragraph in §4 (market-making, immune system, etc.) is the most speculative content. Each named instance has plausible but unverified template-fit; cited only as illustration of the template’s *shape*, not as derived AAT instances. The voice intentionally remains suggestive rather than asserted.
- The deps list includes 16.5 and AT even though they currently sit in Part III Ch.5 and Appendix A respectively. Their content is essentially Part I — multi-dimensional persistence depending only on the chapter’s central result and 3.8 — and 16.5’s own OUTLINE comment flags that it could plausibly live in Part I appendices. Discussing them in the chapter-end is consistent with the math.

Incidental audit gold (2026-05-30 sweep). Cross-audit “wandering thoughts” / §14-ideation lifted from the de-novo auditors’ working dirs (audit-routing-instructions.md §8). Orthogonal pedagogical / framing material staged for a later Brief/Discussion-promotion pass — kept separate from certified theory-fix findings. Thin coverage: this is a chapter-end impl-*** synthesis segment that de-novo auditors were instructed to avoid as priming-heavy, so only one digested reflection reached it (Codex/Claude, AUDIT-WORKING-526815).

3. Follow-up items.

- **Landauer-coefficient conversion (resolved 2026-05-30).** The flagged dissipation floor read “roughly $0.35 n\alpha k_B T$ per unit time”; the 0.35 mixed the bit and nat conventions (a spurious $1/\ln 2 = \frac{1}{2\ln 2} \approx 0.36$ — applied to a nats rate under a per-nat $k_B T$ constant). Verified first-hand that $\dot{R}$ is genuinely a nats rate (Gaussian-OU RDF in natural-log form) and the segment’s own stated convention is $k_B T$ per nat, so $n\alpha/2$ nats/time $\rightarrow \frac{1}{2} n\alpha k_B T$ per time. Corrected to $\frac{1}{2} n\alpha k_B T$ at its canonical home #deriv-persistence-cost and propagated here (Codex/Claude, AUDIT-WORKING-526815).
- **Provenance / status markers for the synthesis.** As a chapter-end synthesis whose proof-burden sits in home segments (several of which are downstream in OUTLINE order), the segment carries heavy first-pass priming; the auditor suggested a prominent “depends on later / appendix home segments” note at the top, and that the speculative cross-domain-transfer paragraph (§4) be unambiguously marked as illustration of the template *shape*, not as derived AAT instances (Codex/Claude, AUDIT-WORKING-526815). The existing Working Notes already state both policies; this is a request to surface them in the body for a reader who does not reach the Working Notes.

5. Candidate figures.

- **Implication hub-and-spoke with provenance markers.** The persistence condition at the center, five implication spokes (information-rate cost / weak-direction bottleneck / structural-vs-parametric failure / template economy / trajectory identity), each spoke carrying a small status marker — *proven earlier* / *future-home-segment dependency* / *speculative transfer* — which communicates the synthesis honestly rather than dressing it as a single theorem (Codex/Claude, AUDIT-WORKING-526815).

2. Candidate Discussion.

- The auditor affirmed the “single template, six instantiations” economy (§5) as a potential *strong positive* once the downstream template is verified to actually support the promised instances — worth foregrounding as a chapter-level contribution if it holds (Codex/Claude, AUDIT-WORKING-526815).

Candidate experiment (2026-08-12, agent-proposed during an incident-response & forensics applications read). The channel-capacity floor’s LLM prediction in §1 (“context-stuffing helps to a point, then degrades — structural rather than a vibe”) is a falsifiable, currently-testable claim: sweep context budget against effective adaptation rate on a drifting task and check for the predicted capacity-limited tempo ceiling $C \geq \mathcal{T}/2$. An empirica registration here would give the framework one of its most externally-legible validation points, since the qualitative phenomenon is already widely observed but lacks a structural account. Agent-proposed forward pointer.

Part II

Agentic Systems: Actuated Adaptation

Scope: agents that not only track reality but aim at something — Part II adds objectives and strategy alongside the reality model. This is the **actuation** half of *Adaptation and Actuation Theory*: Part I develops adaptive systems in general; Part II develops the goal-directed layer built on top.

Headline contribution. Part II’s distinctive results are: the satisfaction gap / control regret split (8.4, 8.5) with the convention hierarchy (6.6) governing the inferential force of each diagnostic from local (C1) through global (C3); the G_t complexity bound (in 10.1); the graph-structure uniqueness argument with CMC-based Markov property (Q); and the survival classification (#result-section-ii-survival) showing that Part II’s definitional and structural architecture transfers exactly to fully-coupled (Class 3: Coupled) agents in 16/24 cases, with the remaining results either approximating with bounded error (5/24) or naming the precise modification required (2/24, plus 1 boundary-defining failure). The Class 3 (Coupled) bias bound underlying the approximate-survival category is now a conditional theorem under named sub-scopes (AB), not order-of-magnitude guidance. The reach is wider than the most-restrictive scope condition suggests; the lattice below names exactly which results live where.

Scope lattice. Part II’s results sit at different points in a stacked lattice of scope conditions, not a single uniform scope. Reading inward from the broadest:

1. **Adaptive scope** (1.5): observation under uncertainty. Part I machinery applies here without further restriction.
2. **Agency scope** (1.6): adaptive + at least binary action choice + at least one action with Pearl-level-2 causal contrast. Part II’s definitional architecture — $X_t = (M_t, G_t)$, $G_t = (O_t, \Sigma_t)$, the agent spectrum — is already non-vacuous here.
3. **Learning-agent scope** (7.2 Scope Narrowing): agency + the agent must acquire or refine its interventional structure during operation. Pre-compiled controllers (PID, LQR, hardcoded reactive policies) sit in agency scope but outside learning-agent scope — their causal mapping was supplied by a designer with external Level 2 access. Most Part II dynamics — the orient cascade, edge-update via gain, strategic calibration, the satisfaction-gap / control-regret diagnostic loop — operate in learning-agent scope.
4. **Class 1 (Separated) architecture**: directed separation (6.3) holds by construction — epistemic processing f_M has no causal path from G_t . Examples: Kalman filter + LQR, modular RL with separate world model, military intelligence separated from operations, tool-use AI agents with separate perception and planning modules. **Class 3 (Coupled) agents** — including transformer-based LLMs where attention processes goals and observations together — lie outside Class 1. Class 2 (Partial) sits between, with coupling quantified by $\kappa_{\text{processing}}$ (6.3).

The lattice stacks. Kalman + LQR is in Class 1 (Separated) but outside learning-agent scope (the control law is pre-computed); a designer applying Part II’s diagnostic core to Kalman + LQR is using only the architecture-survival half, not the learning-dependent dynamics. The orient cascade, in particular, requires both Class 1 (Separated) directed separation (for sequential resolution) and learning-agent scope (for Σ_t revision to be meaningful). The bias-bound theorem and survival classification (AB, #result-section-ii-survival) trace exactly which Part II results require which lattice entries — see those segments for the per-result classification.

Part III’s composite-agent scope sits below this lattice as a fifth tier and inherits the Class-1-(Separated)-sub-agents → Class-2-(Partial)-composite refinement of directed separation (13.3, 16.1).

Part I machinery applies regardless of architecture. Adaptive dynamics on the epistemic substate M_t — mismatch, gain, tempo, persistence — operate independently of how f_M relates to G_t . Class 3 (Coupled) agents are adaptive systems in the Part I sense; what they lose is the clean factorization that gives Part II’s coupled-formulation results their definitional simplicity. The coupled formulation Class 3 (Coupled) agents require is the subject of 03-llm-core/, which inherits Part II’s surviving 16/24 architecture and adds the additional machinery (#scope-observation-ambiguity-modulation, ambiguity-gated bias bounds) that the coupled regime calls for.

“If a man knows not to which port he sails, no wind is favorable.” — Seneca

Chapter 5

Meta-Architecture I: The Certificate and its Boundary, Scope, and Identity Facets

Part II's chapters develop the actuated-agent machinery on top of Part I's adaptive substrate. Before walking them, this chapter introduces the cross-cutting vocabulary that the five chapters that follow — and the parts beyond — surface as instances. Three of the four facets of the stability-certificate spine operationalize within Part II; the fourth (dynamics on architectural class) develops once composite-agent machinery is in hand and is introduced at the opening of Part III.

The certificate and its facets. The framework's analytical surface is organized around the **stability-certificate spine** (5), grounded in the Lyapunov / converse-Lyapunov equivalence: operator-sector contraction in some metric and exponential stability are interchangeable witnesses for the same property. The certificate has facets that AAT's results inhabit; the facets are provably distinct (Helmholtz / Sylvester / Mori-Zwanzig), and the plurality is the content. Three facets operationalize in Part II: **M1 — Identifiability floor** (5.1, boundary facet; structural no-gos on what can be inferred from limited data, four instances of which three form a Sylvester rank-collapse subclass irreducible by Sylvester's law of inertia); **M2 — Separability pattern** (5.4, scope-of-existence facet; the separable-core / structured-repair / general-open trichotomy across eight ladders); **M3 — Additive coordinate forcing** (5.5, forced-identity facet; uniqueness theorems on AAT-internally-motivated axioms at four representational layers).

M1's two sister meta-segments at the same boundary facet. Two further meta-segments sit alongside M1, agent-side and designer-side. 5.2 (agent-side): the two-routes-exhausts cluster around #form-objective-functional's single-interface commitment — the agent-side adaptive-substrate route (Result G' via #deriv-self-actuation-grounding) and the principal-side protocol-commitment route (Cohen-2022 strengthened via #deriv-reward-channel-learning-no-go) together exhaust the structural complement of V_{O_i} 's information narrowness. 5.3 (designer-side): AAT-side translations of three classical mechanism-design no-gos (Gibbard-Satterthwaite 1973-75, Myerson-Satterthwaite 1983, Arrow 1951); designer-with-construction-task no-gos that contrast with M1's agent-with-inferential-task no-gos via remedy (design-constraint relaxation vs information augmentation). The suffix discipline (*-floor / -impossibility*) encodes the actor-positioning distinction at slug-grep depth.

A style claim atop the boundary facet. 5.6 names the three-move discipline — name the floor (Pearl-Bareinboim CHT / Cramér-Rao / Liberzon common-Lyapunov nonexistence / Khasminskii recurrence / sharp algebraic supremum), name the escape (unique broadly-available framework machinery), treat the no-go as apparatus (elevating escape from "useful" to structurally required; when the no-go *is* the result, on the critical path) — that the M1-cluster and M4-cluster instances run. It is a *style* claim about how AAT states certain results, sitting atop the boundary facet rather than as a fifth facet; the catalog spans the three M1-cluster clusters and connects forward to the M4-cluster instances introduced at the opening of Part III.

A second style claim sits alongside it. 5.7 names the **anti-collapse** discipline — AAT’s habit of refusing to merge two things a naive model would treat as one, because the collapse hides a difference that routes to a *different repair* (the deeper principle: individuate causes and effects at the grain the remedy cares about). It is planted at its first instance (β -vs- ρ , 2.2, Part I), developed here as cross-cutting vocabulary, and recurs across all three Parts. Like the constructive-impossibility posture, it is an epistemic-architectural recognition — how the framework states results — not an additional facet of the certificate.

AAT’s cross-sectional structure is the geometry of a single object — the **equilibrium stability certificate**, the positive-definite form whose existence certifies that an agent can correct itself faster than its world drifts; operator-sector is that object’s interior, the separability pattern its scope of existence, additive-coordinate-forcing its forced identity, and the identifiability floor its boundary, with composition the question of whether it survives projection.

The three meta-segments 5.4, 5.1, and 5.5 read, separately, as three independent organizing insights that happen to recur. This segment names the object they are facets of. The relationship is the same one 5.5 already runs at smaller scale (“layer-specific manifestations of a single geometric object”), raised to the framework: not a fourth meta-pattern *alongside* the three, but the spine the three are projections of.

For an agent with error dynamics $\dot{e} = -F(e)$ about an equilibrium e^* ($F(e^*) = 0$, $F \in C^1$ near e^* , Jacobian $J := DF(e^*)$), a **stability certificate** is a symmetric positive-definite $\mathcal{M}$ for which the one-point sector condition holds in the $\mathcal{M}$ -inner-product on a ball $\mathcal{B}_R(e^*)$:

Definition (stability-certificate).

$$\langle F(e), e - e^* \rangle_{\mathcal{M}} \geq \kappa \|e - e^*\|_{\mathcal{M}}^2, \quad \kappa > 0. \quad (\text{C})$$

The certificate is not unique: it is whatever positive-definite form makes the dynamics contract. In the recurring sub-cases it specializes — to the Fisher information for Bayesian agents, to $(P^-)^{-1}$ for Kalman agents, to the loss Hessian for gradient agents, and to a plant-selected Lyapunov metric for linear-Hurwitz or PID agents. These are not four separate stories; they are one object under four certificates.

The object is load-bearing only because its existence is not a definition but an equivalence: **a stability certificate exists iff the agent is exponentially stable about its target** — operator-sector in *some* inner product and exponential stability are the same statement, with the certificate as the converse-Lyapunov witness, and the certificate admits a strict strength ladder $R0\text{-loss} \Leftarrow R0\text{-strict} \Leftarrow R1 \Leftarrow R2$ (widest Lyapunov-stable; strict-contractive one-point/local; cocoercive; Čencov-forced). This is the segment-level form of the contraction-over-drift organizing principle. The equivalence, the ladder, and the proof are stated and derived exactly in H; this spine cites that result and builds the cross-sectional reading on it rather than re-deriving it.

Result (cited: #result-certificate-existence).

The certificate is one object; the cross-sectional meta-patterns are its facets on the positive-semidefinite cone $\mathbb{S}_{\geq 0}^n$:

Discussion.

Table 5.1

<i>Facet</i>	<i>Meta-segment</i>	<i>What the facet is</i>	<i>Canonical home</i>
Interior	G, AI	$\mathcal{M} > 0$ on the scope ball, Helmholtz S/A split: $S_{\mathcal{M}} > 0$ gives R0-strict contraction; $S_{\mathcal{M}} = 0$ at $\mathcal{M} = I$ gives R0-loss conservation; $S_{\mathcal{M}} \geq 0$ with nontrivial kernel is the Conley-decomposed mixed case	the template segments
Scope of existence	5.4	the region where a certificate exists at all (separable core / structured repair / general open)	M2
Forced identity	5.5	<i>which</i> certificate: Čencov forces $\mathcal{M} = \text{Fisher}$ uniquely in statistical scope; matched (existence-only) elsewhere	M3
Boundary	5.1	$\mathcal{M}$ drops rank ($\partial \mathbb{S}_{\geq 0}^n$): the inferential task is structurally impossible	M1
Projection behaviour	13.1	whether a <i>common</i> certificate survives coarse-graining; the closure defect ε^* is the certificate's projection-residue	composition-closure

Each meta-segment retains its own canonical home and per-instance derivations; this segment claims only the recognition that they are facets of one object, and what that buys (Discussion below).

A tempting reading is that the certificate's failures are one obstruction seen three ways (a single "failure of integrability"). They are not. The three failure modes are irreducibly distinct, each invariant under the others' degrees of freedom:

Discussion.

- **Forced-identity failure — Helmholtz–Hodge.** J non-symmetric $\Rightarrow$ the field is not a gradient $\Rightarrow$ no potential $\Rightarrow$ the certificate is *matched* (converse-Lyapunov existence), not *forced* (Čencov). A non-symmetric Hurwitz J still has a certificate (it is *not* on the boundary), so this is an M3 failure, not an M1 one. Invariant: symmetry of J .
- **Existence failure — Sylvester's law of inertia.** The certificate drops rank. Every coordinate/metric change acts on the certificate by congruence; congruence preserves inertia; so a rank-deficient certificate is rank-deficient in *every* coordinate. The boundary is invariant under the agent's entire representational freedom (that freedom *is* the congruence orbit); the only escape is rank-augmentation — genuinely new information, not a re-mapping. (Detailed in 5.1's Sylvester finding.) Invariant: inertia under congruence.
- **Projection failure — Mori–Zwanzig / Schur.** Coarse-graining is a non-invertible projection. The certificate-as-metric survives (the Schur complement of a positive-definite form is positive-definite) but the *dynamic* guarantee does not: the closure defect ε^* equals the norm of the Mori–Zwanzig memory commutator, zero exactly when the resolved subspace is J -invariant. Invariant: J -invariance of the resolved subspace.

Each obstruction is untouched by the others' freedoms: a metric change does not fix non-invariance; projection does not fix non-symmetry; rank-augmentation does not fix a memory kernel. That mutual invariance is the reason the cross-sectional structure is *several* meta-patterns and not one — stated as a structural fact rather than left as "they are different concerns."

Discussion-grade at the organizing-principle level. This segment names a single object behind separately-derived results; it is not itself a new theorem. What is derivative here is the recognition that the cross-sectional meta-patterns are facets of one certificate — the recognition, not a fresh derivation, is the content.

Constituent results retain their own, higher status. The certificate-existence equivalence (Lyapunov theorem; Formal Expression “anchor”) is *exact* at the linearized level and *exact-with-standard-remainder* locally. The certificate-strength ladder R0-loss/R0-strict/R1/R2 is *exact* as an ordering of conditions (R0-loss’s pure case via Letcher 2019 Helmholtz characterization; mixed case via Conley 1978 universal decomposition). The boundary-irreducibility mechanism (Sylvester’s law) is *exact*; its full statement and per-instance verification live in 5.1. The projection-residue identification (ϵ^* = memory-commutator norm) is *robust-qualitative*, with the per-case content in 13.1. The forced-vs-matched distinction (Čencov statistical-scope-only) is established in 5.5.

Scope honesty. The anchor equivalence is linearized/local — this is the level at which AAT’s persistence results already operate (sector conditions, contraction templates, the bridge lemma all linearize about the equilibrium), so it is not a weakening relative to the rest of the theory, but it is a genuine scope statement and is not papered over: there is no claim that one-point operator-sector is equivalent to *global* exponential stability (it is not, in general). The synthesis claim — that AAT’s *entire* cross-sectional structure is this one cone — is *robust-qualitative*: each per-facet identification is exact or cited, and the “all of it” is as strong as those identifications jointly, no stronger. Whether exactly three obstructions exhaust the failure modes is not proved; three are established and each is exact, but exhaustiveness is open.

Max attainable: *discussion-grade* for the organizing claim (it is a presentational spine, not a derivation). The anchor equivalence and the Sylvester mechanism retain their *exact* status at their own canonical homes.

What the spine buys. Three things. (i) It converts “AAT has several recurring organizing insights” into “AAT’s cross-section is one object’s facets,” which is the difference between a catalog and a structure — a reader who holds the certificate picture can predict where the meta-patterns will bite (interior / scope / forced-identity / boundary / projection) rather than meeting them as separate surprises. (ii) It grounds the long-standing organizing slogan — *an adaptive system is an operator whose contraction rate exceeds its disturbance rate* — at segment level, as the anchor equivalence rather than as a heuristic. (iii) It makes the framework’s scope-honesty sharper: the prior positioning that the identifiability floor is “orthogonal” to the contraction machinery is replaced by the exact geometric statement that the floor is the *boundary* of the very cone whose interior the contraction machinery is, with the boundary held invariant by Sylvester’s law against the framework’s only representational freedom.

Relationship to the facet segments. This segment cross-references; it does not restate. 5.1 remains the canonical home of the floor instances and the Sylvester mechanism; 5.4 of the separable-core / structured-repair / general-open ladders; 5.5 of the Čencov/Cauchy-FE forcing; G and AI of the contraction machinery; 13.1 of the closure defect. The spine’s contribution is the cross-segment recognition and the anchor equivalence, not any per-facet content.

Why the plurality of obstructions is a feature. Had the three failure modes collapsed to one, the spine would be a single clean theorem — elegant but false on the evidence (the integrability collapse fails: a non-symmetric Hurwitz field has a certificate yet no potential; congruence is invertible while projection is not). The honest structure is more useful: it tells a future agent precisely *not* to seek a single mechanism unifying the floor, the closure defect, and the forcing failure — they are Sylvester, Mori–Zwanzig, and Helmholtz respectively, and the proof of their distinctness is their mutual invariance. The unification is real at the object level and the no-go is real and plural at the failure level; both are load-bearing.

Complementarity with the existing meta-segment framing. 5.4 names AAT’s positive half and 5.1 its negative half; 5.5 names the constructive half. The spine names what all three are halves of. The three remain the right reading lenses for any individual segment; the spine is the reason those lenses compose rather than merely coexist. 5.6 sits atop the boundary facet specifically — naming the framework’s cross-instance *style* (name the floor, name the unique broadly-available escape, treat the no-go as load-bearing apparatus) rather than introducing a fifth peer cross-sectional meta-pattern.

The Helmholtz S/A split of the linearized certificate is the spectral/representation dual of the Interior facet, not a fifth meta-segment. The Interior facet (G, AI) is the *static* statement “ $\mathcal{M} > 0$ on the scope ball: the contraction holds.”

Read through the Helmholtz decomposition of the linearized certificate, $S_{\mathcal{M}} := \frac{1}{2}\mathcal{M}^{-1}(\mathcal{M}J + J^{\top}\mathcal{M})$, the facet decomposes into two interior sub-regions: $S_{\mathcal{M}} > 0$ gives R0-strict contraction (the prior Interior content, $\dot{V} \leq -2\kappa V$ with $\kappa > 0$); $S_{\mathcal{M}} = 0$ at $\mathcal{M} = I$ gives the antisymmetric-subclass R0-loss case (certificate exists at $\sup \kappa = 0$, V exactly conserved). The mixed case ($S_{\mathcal{M}} \geq 0$ with nontrivial kernel) is Conley 1978's Fundamental Theorem of Dynamical Systems localized at the linearization — chain-recurrent subspace + strongly-gradient-like complement, with the full four-case spectral partition and the non-antisymmetric R0-loss counterexample at $\mathcal{M} \neq I$ worked in H Derivation. Run through this segment's own test ("is a candidate organizing pattern a new facet, or a genuinely new object?") returns **neither**: it is new internal structure on the existing Interior facet — the spectral/representation dual of contraction-strength, splitting the interior into strict (contracting) and loss (Lyapunov-stable-only) sub-regions — not a peer fifth facet. The four-facet structure stays four; the cone gains no sixth facet and the framework no fifth meta-segment. The theorem-anchors are Letcher-Balduzzi et al. 2019 Lemma 1 + Definition 2 (pure-case sufficient construction; the Hamiltonian-game vocabulary saturating the antisymmetric sub-class) and Conley 1978 §8.1 (mixed-case universality); the worked composite instance — joint best-response field at saddle-only Nash in cyclic-distributional games (Cheung-Piliouras-Tao 2021 Theorem 19) — lives on the $R1 \rightarrow R2$ boundary of the composite dynamic-regime axis (AF), where $R2$'s joint Jacobian linearizes to the antisymmetric R0-loss-at- $\mathcal{M} = I$ sub-case. The structure becomes load-bearing at the $R0\text{-strict} \leftrightarrow R0\text{-loss}$ boundary — strict contraction degrading toward conservation under loss of $S_{\mathcal{M}} > 0$ at the chosen metric — and the *cross-row recognition* follows: rows 03 + 10 + 14 of ref/prior-art-analysis/ (sector condition / persistence template / strategic composition under goal divergence) share the Interior-via-Helmholtz reading at the single-agent linearization level; row 17 (credit assignment) lives on the M1 boundary facet via Sylvester rank-collapse, *not* on the Interior facet — the cross-row unification is scoped to rows 03/10/14 only.

The accumulation-typing pattern is the temporal dual of the Interior facet, not a fifth meta-segment (the D3 register answer). The Interior facet (G) is the *static* statement " $\mathcal{M} > 0$ on the scope ball: the contraction holds." Read in *accumulation/temporal* vocabulary it is the statement "the operator carrying a per-step residue to its accumulated consequence is bounded" — and, per that template's typed-bridge framing, this is **two-model, not one operator** (Model D / Model S; the $1/\alpha$ vs $1/\sqrt{\alpha}$ scaling and the Cor-A.IS.1 categorical containment dichotomy are the fingerprint that they are different functionals, not one read in two norms). The recurring *accumulation-type confound* — asking about a per-step residue in the vocabulary of its accumulation, which dissolved the $\epsilon^*(N)$ poly-vs-exponential question and recurred across the framework — run through this segment's own test ("is a candidate organizing pattern a new facet, or a genuinely new object?") returns **neither**: it is a new *reading* of an existing facet (the temporal/representation dual of the Interior), so the cone gains no sixth facet and the framework no fifth independent meta-segment. The notation that carries it is durable canon (NOTATION.md §"Accumulation typing"); its theorem-anchor is the Interior facet's template. The structure becomes load-bearing exactly at the *Interior*→*Boundary transition* — contraction degrading toward rank-loss — where two operator families separate: $\mathcal{A}_{\text{ref}}$ (a reflected Lindley/Loynes walk whose content is the driftless $\mu = 0$ boundary) instantiates the template on the turnover index (#der-identity-continuity-threshold), while $\mathcal{A}_D$ (the linear destroy-and-reconstruct contraction whose norm diverges as the gap closes) is the honest *non-transfer* boundary (#der-turnover-information-recursion). The $\mu = 0$ boundary of the one is the pole of the other, which is why they are distinct operators and treating one as a linearization or normalization of the other is a category error. This is the framework's worked answer to its own "new pattern?" test from the time side: *no* — *it is this existing object seen from the time side*, complementary to the Helmholtz S/A paragraph above which gives the spectral-side worked answer.

Brief: Think of an adaptive agent as trying to stay on a moving target, and ask one question: is there a way of measuring "how far off am I?" such that every correction the agent makes provably shrinks that measure faster than the world pushes it back out? That measuring-stick is the stability certificate. The whole cross-sectional skeleton of AAT turns out to be facts about *this one stick*: the agent can keep up exactly when the stick exists and is positive (operator-sector — the interior); the framework's reach is exactly the region where some such stick exists (the separability pattern); when the stick is pinned down uniquely it is pinned by one classical invariance theorem and only in the statistical case (additive-coordinate-forcing); and the agent provably *cannot* keep up exactly when the stick goes flat in some direction (the identifiability floor — the boundary), with a second classical theorem (Sylvester's law of inertia) proving no change of

measuring units ever un-flattens it — only genuinely new information does. A thoughtful non-specialist can carry the whole structure away from the one picture: existence of the stick = can adapt; flatness of the stick = a blind spot no re-measurement fixes; uniqueness of the stick = a special (statistical) privilege, not the general case; and looking at the stick through a coarse lens (composition) keeps its shape but loses its guarantee by exactly a memory term.

Impact: Reorganizes AAT’s self-description from three independent meta-patterns plus a contraction mechanism into one object with four facets and one anchor equivalence, so the framework’s cross-section can be read predictively rather than as a catalog of separate recurrences. Grounds the organizing slogan (contraction-rate-exceeds-drift) at segment level via the Lyapunov-theorem equivalence, discharging the long-standing “not yet surfaced at segment level” status. Sharpens the scope-honesty posture: “the floor is orthogonal to operator-sector” becomes “the floor is the boundary of the cone whose interior is operator-sector, held invariant by Sylvester’s law against the framework’s only representational freedom.” Bounds its own claim honestly: the unification is at the object level; the three failure obstructions are provably *distinct* (Helmholtz / Sylvester / Mori–Zwanzig), and that plurality is precisely why AAT carries multiple cross-sectional meta-patterns rather than one — now a stated structural fact instead of an intuition. Gives every future organizing-pattern candidate a test: is it a new facet of the certificate, or a genuinely new object?

Novelty Claim: *Claim recognition* that AAT’s cross-sectional meta-patterns (separability, identifiability-floor, additive-coordinate-forcing) and its contraction machinery are facets — interior, scope-of-existence, forced-identity, boundary, projection-behaviour — of a single object, the equilibrium stability certificate; together with *claim synthesis* binding the Lyapunov-theorem certificate-existence equivalence, the Sylvester-law boundary-irreducibility, and the Mori–Zwanzig projection-residue into one cross-sectional structure. The constituent theorems are classical; the contribution is the recognition that AAT’s separately-derived meta-patterns are one object’s facets and that their failure modes are provably plural.

Related Work: - Lyapunov, A. M. (1892), *The General Problem of the Stability of Motion*; Khalil, H. K. (2002), *Nonlinear Systems* 3rd ed., Thm 4.6 (found 2026-05-14) — *formal antecedent* — the certificate-existence equivalence (operator-sector in some metric $\Leftrightarrow$ Hurwitz). - Sylvester, J. J. (1852), *Phil. Mag.* 4(23):138–142; Horn & Johnson (2013), *Matrix Analysis* 2nd ed., Thm 4.5.8 (found 2026-05-14) — *formal antecedent* — the boundary-irreducibility mechanism; full treatment in 5.1. - Mori (1965) / Zwanzig (1961); Chorin, Hald & Kupferman (2002), *Physica D* 166:239 (found 2026-05-14) — *formal antecedent* — the projection-residue (memory kernel) underlying the composition facet; per-case content in 13.1. - The facet segments 5.1, 5.4, 5.5 — *adjacent* — each carries its own per-instance prior-art landscape; the spine adds the cross-segment object, not the per-facet priors.

Search Log: - 2026-05-14 (*targeted*): The recognition was assembled from classical pieces (Lyapunov / Sylvester / Mori–Zwanzig) plus the framework’s own meta-segments. The search target was whether “an integrated agent theory’s cross-sectional meta-patterns are facets of a single stability-certificate cone” appears as an articulated structure elsewhere. Not found at this depth; the constituent theorems are textbook and the per-facet identifications are individually well-precedented, but the cross-segment unification as a framework spine is a fresh presentational recognition. Expected to remain *recognition/synthesis*-tier under deeper search — the pieces are classical; the assembly is the contribution. Per-facet comprehensiveness is inherited from the facet segments, not from a fresh cross-facet search.

- **Provenance.** The certificate-spine recognition, the anchor equivalence, the Sylvester boundary mechanism, and the broken-integrability-triad result were worked out in the 2026-05-14 operator-family-unification cycle (the deep push of the predecessor C1 “2-instance-plus-1-consequence” question and the prior co-owner “do not elevate unless O-BP10 surfaces at segment level” gate, now met by the anchor equivalence). See CHANGELOG 2026-05-14 for the cycle narrative. The originating spike is absorbed archaeology, not a live reference; the load-bearing content is in this segment, H, and the 5.1 Sylvester finding.
- **Interior-facet Helmholtz S/A refinement (2026-05-22).** The Interior facet’s two-sub-region structure (R0-strict / R0-loss split via Helmholtz S_M) landed via Track CR Phase 4 per spikes/spike-integration-reconciliation-2026-05-22/99-verdict.md §5 Phase 4 + spikes/spike-enrichment-cluster1-2026-05-21/02-r0-loss-derivation.md

(§05 fifth-facet test). The refinement parallels the line-84 accumulation-typing paragraph (“new facet or new object?” returns “neither — new internal structure on an existing facet”). Four-facet structure preserved; the L4 cross-row recognition (rows 03+10+14 share the Interior-via-Helmholtz reading; row 17 lives on M1 boundary via Sylvester) is anchored here at the spine rather than scattered across the four rows. The full R0-loss characterization (the four-case spectral partition, including the non-antisymmetric case-(iii) counterexample) lives in H Derivation, not duplicated here.

- **Dependency rationale (for Gate-1 audit).** depends: lists result-certificate-existence (the anchor equivalence this spine builds the cross-sectional reading on — a genuine dependency, consumed not merely recognized) plus result-sector-persistence-template and deriv-sector-condition (the persistence machinery the anchor’s drift half rests on). 5.1, 5.4, 5.5, and 13.1 are **cross-referenced as facets, not depended on**: per FORMAT.md Gate 1, a dependency is genuine only when the segment uses the referenced segment’s definitions/results, not when it is recognized or related in Discussion. The spine *recognizes* the meta-segments as facets of one object; it does not consume their definitions to make its claim. This is why the spine is correctly placed *before the three meta-segments* in OUTLINE (it is the object they are facets of, so it reads first) without an ordering violation — the facet relationship is lateral recognition, not a dependency edge — while being placed *after* #result-certificate-existence (a genuine dependency, which therefore precedes it). Treat the facet #... references as expected forward/lateral cross-refs (FORMAT.md §Cross-References).
- **Provisional slug.** disc-stability-certificate. Alternative considered: disc-certificate-cone (names the geometry rather than the object). Subject-noun discipline favours the object (“the stability certificate”); the cone/interior/boundary is what the segment *says about* it. Route through the naming pipeline if a better name surfaces; the spike verdict floated both.
- **Findings Brief sharpening queued (2026-05-28; audit 384279 §D Hyp 5).** The current Findings Brief (line 92) uses a “measuring-stick” analog for the certificate that conflates two things slightly — the metric (the yardstick) and the contraction property (every correction shrinks the measure). A fresh reader could interpret “measuring-stick” as just the metric without grasping the load-bearing part — that the metric is *chosen so* contractions provably shrink it. The Walton-bathtub analog in #persistence-and-limits-intro has cleaner one-to-one mapping. Two sharpening options queued at spikes/PROPOSED.md Tier 1 (2026-05-28 row): (i) tighten the Brief to one mechanism (correction-faster-than-disturbance with the certificate as the positive-definite form making that contraction provable); (ii) extend the bathtub-parallel — bathtub with graduated marks where the certificate exists exactly when there are marks calibrated such that every drop entering provably tips the gauge downward faster than inflow. Option (ii) recommended because it preserves continuity with the established bathtub analog. Editorial work, ~30 minutes; load-bearing for Feynman-criterion isomorphism in the spine’s pedagogical surface.
- **Provisional OUTLINE position.** Placed in ## *Appendices* Details immediately before 5.1, as the lead of the four-row meta-segment cluster (spine → M1 → M2 → M3). Provisional because: (a) the meta-segments may eventually warrant their own chapter rather than Appendix-A residence; (b) if the OUTLINE preamble is reframed to lead with the spine, the cluster likely moves to a more prominent position. Both are propagation steps below, not landed here.
- **Propagation plan (ordered by commitment; steps 6–7 are framework-voice keystones gated on Joseph, not auto-executed):**
 1. 5.1 — cross-ref line in its “complementarity” Discussion paragraph naming the spine as the object whose boundary it is. (Sylvester finding + Working Note already point here.) Cross-ref only.
 2. 5.4 — parallel line: it is the *scope-of-existence* facet (where the certificate > 0). Cross-ref only.
 3. 5.5 — frame its (PI)/Čencov content as the *forced-identity* facet. Cross-ref only; Čencov machinery’s canonical home stays M3.

4. **G / AI** — one Discussion sentence: the template condition is the certificate interior; cross-ref the spine for the cone reading. No formal change.
 5. **13.1** — Discussion line framing ϵ^* as the certificate's projection-residue and Liberzon as “no common certificate”; cross-ref the spine. (Its Mori–Zwanzig Working Note already exists.)
 6. **O-BP10 keystone (Joseph's call)**. Recommend this segment *is* O-BP10's segment-level home: the slogan is its one-sentence summary, the equivalence is its anchor. The PROPOSALS Bundle-1 O-BP10 entry then points here. Do not auto-rewrite Bundle 1.
 7. **OUTLINE preamble reframe (highest commitment; Joseph's call)**. OUTLINE.md line 17 currently opens “Three meta-segments form AAT's cross-sectional structure: 5.4 ... 5.1 ... 5.5 ...”. Proposed replacement, for Joseph's confirmation before it goes live in the auditor-visible preamble: *“AAT's cross-sectional structure is one object — the equilibrium stability certificate (5). Its interior is operator-sector (the contraction machinery); its scope of existence is the separability pattern (5.4, positive half); its forced identity is additive-coordinate-forcing (5.5, constructive half, Čencov-forced in statistical scope only); its boundary is the identifiability floor (5.1, negative half); and its behaviour under coarse-graining is the composition-closure defect. Reading any segment through the certificate and its facets surfaces what makes it load-bearing: whether a certificate exists for it, which one is forced, what boundary it abuts, and whether it survives projection.”* Execute only on Joseph's confirmation, having seen the segment land first.
- **Open edges (from 99-verdict.md)**. Anchor equivalence is linearized/local (stated in Epistemic Status). “Exactly three obstructions” is robust-qualitative, not proved exhaustive — a fourth (e.g. non-autonomous certificate drift for time-varying systems) is not searched. Sylvester is proved finite-dimensional; the function-space (M_t for logogenic agents) extension is unchecked and flagged for any future logogenic application, not load-bearing for the AAT-core claim.

Orthogonal pedagogical / framing material harvested from de-novo auditors' working-dir reflections, kept separate from certified findings. **Coverage:** reached as a named segment by audit 384279 (note 33), the one auditor who walked the current Meta-Architecture-I OUTLINE with the spine first; most others predate the spine's landing. Preserve the finding-vs-framing conflation.

- **The Findings-Brief measuring-stick analog conflates two things** — the certificate-as-metric (the yardstick) and the contraction property (every correction provably shrinks it). A fresh reader could read “measuring-stick” as just the metric, missing the load-bearing part: that the metric is *chosen so* contractions provably shrink it. Two sharpening options: (i) lead with one mechanism — “the positive-definite form that makes that contraction provable”; (ii) extend the established Walton-bathtub analog — “a bathtub with graduated marks calibrated such that every drop entering provably tips the gauge downward faster than the inflow” — recommended for continuity with the Ch.4 chapter intro (Claude, 384279). (*Already captured in this segment's Working Notes “Findings Brief sharpening queued” bullet and at spikes/PROPOSED.md Tier 1; recorded here for completeness so the gold-lift trail is whole.*)
- **Catalog-to-structure as the spine's plain-language payoff**. The three-obstructions-are-distinct argument (Helmholtz / Sylvester / Mori-Zwanzig) “is the kind of *structural* recognition that converts a catalog of insights into a theory” — and “the plurality is the content” is the right read. A candidate one-line framing of *why the spine earns its keep*: it lets a reader predict where the meta-patterns will bite (interior / scope / forced-identity / boundary / projection) rather than meet them as separate surprises (Claude, 384279, echoing the segment's own “What the spine buys” §(i)).

- **The accumulation-typing temporal-dual paragraph reads as a heavy lift in a Discussion section.** An auditor flagged the line-86 paragraph as “*very* dense,” asking the reader to grasp two new operators ($\mathcal{A}_{\text{refl}}$ Lindley/Loynes vs $\mathcal{A}_D$ destroy-and-reconstruct) and their separation at the $\mu = 0$ boundary, and forward-referencing segments not in the auditor’s OUTLINE scan (#der-identity-continuity-threshold, #der-turnover-information-recursion). Suggestion: could plausibly move to WorkingNotes or a sister segment to keep the spine’s Discussion readable (Claude, 384279). A placement / pedagogy follow-up, not a correctness finding.
- **Process-discipline noted as exemplary, worth surfacing in framing prose.** The Working-Notes propagation plan that explicitly gates the two framework-voice keystones (O-BP10, OUTLINE-preamble reframe) on Joseph’s confirmation, and separates auto-executable cross-ref steps from Joseph-gated ones, was named “exemplary cross-cycle handoff documentation” and “exactly the right shape” for peer collaboration (Claude, 384279). Not reader-facing content per se, but a signal that the spine’s provenance discipline is itself legible to a fresh reader.

5.1 Discussion: The Identifiability Floor — A Class of Structural No-Go Results

AAT recurrently applies a **constructive-impossibility methodology** with the following five-element shape:

1. **Setting.** An AAT inferential task — detect a structural property, identify a parameter, distinguish two model classes, recover an internal architecture — restricted to a stated information regime (purely observational data, single-channel observation, marginal-only access, on-policy summaries).
2. **External theorem.** An information-theoretic limit independent of AAT: the Pearl-Bareinboim causal hierarchy theorem (Bareinboim, Correa, Ibeling & Icard 2022); the Cramér-Rao bound on Fisher information; Liberzon’s common-Lyapunov-nonexistence for switched-system contraction; Kalman-Ho canonical-form non-uniqueness for state-space realizations.
3. **No-go.** The external theorem proves that the inferential task is impossible under the regime — no statistic, no Bayesian comparison, no online estimator can succeed using only the available information. AAT does not derive the impossibility; it recognizes the AAT setting falls within the external theorem’s scope.
4. **Boundary characterization.** The conditions under which the regime fails — i.e., the agent has *more* information than the regime allows — admit (partial) identification. Each boundary route maps onto specific AAT machinery already required for other reasons (loop-interventional access, multi-channel observability, observable latents, matched-Tier composite-extended observation, white-box architecture instrumentation).
5. **Strengthened consequence.** The floor elevates the load-bearing role of whichever AAT machinery is the unique broadly-available violation of the regime — from “useful additional tool” to “structurally required by the theory.” Without that machinery the no-go forbids the inference entirely.

The methodology is *not* a negative posture. It is constructive: each floor states “here is precisely what cannot be inferred from limited data; here is exactly which additional capability is required to recover identification; the AAT machinery that supplies that capability is therefore load-bearing in the strongest possible sense.”

Why this discipline matters. Four motivations sit behind the methodology, all visible across the four instances developed below.

The floors strengthen AAT’s foundational machinery. Each instance identifies a piece of AAT machinery as load-bearing in the strongest sense — without it, the corresponding inferential task is *impossible*, not merely *harder*. The loop’s interventional access is not just useful; it is the unique broadly-available violation of the on-policy detection no-go.

Observability of latents is not just convenient; it is the unique route from refuted-by-Cramér-Rao to globally-derived for L1' transfer.

The floors map AAT's limits honestly. The methodology precisely characterizes what AAT cannot do without additional capability. This is scope-marking — at each floor, the theory states what is impossible, what is required to escape the impossibility, and what AAT itself supplies or does not supply.

The methodology gives a unifying template for future no-go results. Adjacent floors not yet formalized (causal-IB extension, misspecification cost, tier-switching cost) are open research directions the shape predicts. Each, if landed, would follow the same five-element trajectory.

The methodology connects AAT to information-theoretic foundations. The external theorems (Pearl-Bareinboim hierarchy; Cramér-Rao bound; Liberzon common-Lyapunov-nonexistence; Kalman-Ho non-uniqueness) are the mature literature AAT inherits. Each instance positions AAT as a consequential application of an established theorem — not a re-derivation.

Position in the certificate spine. This segment is the *boundary facet* of the stability certificate (5). The three rank-collapse instances (Instances 1, 2, 4) are the certificate dropping rank — the boundary of the positive-definite cone whose interior is operator-sector — with the Sylvester-law irreducibility (developed in §"Discussion" below) stating that the framework's representational freedom is a congruence orbit that cannot cross that boundary. Instance 3 is the projection/closure member — a structurally distinct Schur-complement obstruction; the *plurality* of mechanism is itself part of what the methodology carries, and is the reason a single theorem cannot replace the five-element shape. The cross-instance *style* — name the floor, name the unique broadly-available escape, treat the no-go as load-bearing apparatus — is named separately at 5.6 as a posture atop this facet rather than a fifth facet of the spine.

The four current instances follow. Each demonstrates one route through the five-element methodology — different external theorems, different inferential tasks, different boundary routes, different load-bearing machinery elevated. Together they exhibit both the unity of the methodology and the plurality of its applications.

Instances.

Instance 1 — On-Policy L0 Insufficiency Detection (9.2). **Setting.** Detect whether an L0 strategy DAG is causally insufficient (a latent common cause is acting on multiple sibling action propositions) using only the agent's on-policy observation history under sequential short-circuit AND/OR execution.

External theorem. Bareinboim, Correa, Ibeling & Icard (2022) Causal Hierarchy Theorem: there exist SCMs that agree on Level 1 (associational) data but disagree on Level 2 (interventional) questions. Therefore Level 2 distinctions are not in general identifiable from Level 1 data.

No-go. For any L1 world $\mathcal{W}_{L1}$ with latent common cause C , there exists an L0 world $\mathcal{W}_{L0}^*$ with edge probabilities matched to the on-policy regime conditionals such that the on-policy observation distribution is *identical*: $\mathbb{P}_{\pi_{L0}}^{\text{obs}}[\mathcal{W}_{L1}] = \mathbb{P}_{\pi_{L0}}^{\text{obs}}[\mathcal{W}_{L0}^*]$. The L0/L1 distinction is a Level 2 question (it concerns $P(A_2 \mid \text{do}(\neg A_1))$ vs $P(A_2 \mid \neg A_1)$); on-policy data is Level 1; the CHT forbids distinguishing them.

Boundary characterization. Five routes (cf. 9.2): - (a) ε -exploration violates pure on-policy execution $\rightarrow$ partial detection at scale $O(\varepsilon)$. - (b) Joint sibling observability under exploration violates short-circuit censoring $\rightarrow$ strong detection via the pairwise covariance test under 7.3. - (c) Intermediate-state observability at finer grain $\rightarrow$ very strong detection where available. - (d) Structural priors positing common causes $\rightarrow$ prior-quality-dependent. - (e) Direct intervention on the candidate latent $\rightarrow$ strongest where available.

Strengthened consequence. The covariance test under joint observability (route (b)) becomes the unique broadly-available detection mechanism. This sharpens 7.3 from "useful machinery" to "structurally required to escape the no-go." The orient cascade's step 4c (causal-sufficiency check) is no longer "one possible diagnostic" but "the unique broadly-available diagnostic given the structural impossibility of purely on-policy detection."

Tier. *Exact* for shallow strict-prerequisite cases (2-sibling OR or AND with binary common cause). *Robust qualitative* for general DAG topology, soft facilitators, and deeper structures.

Instance 2 — L1' Mixture Identifiability from Single-Channel Observations (R Prop B.7). **Setting.** Identify the mixture parameters $(\theta_C, p_{j|C}, p_{j|\neg C})$ of a soft-facilitator L1' DAG using single-channel observations y_j of one child at a time, with C unobservable.

External theorem. The Cramér-Rao bound (Cramér 1946, *Mathematical Methods of Statistics*, Princeton University Press): the variance of any unbiased estimator is at least the inverse of the Fisher information matrix; if the Fisher matrix is rank-deficient, the bound is infinite in the null directions and no unbiased estimator achieves finite variance there.

No-go. Computing the Fisher information of the mixture model $\mu_j = \theta_C \theta_{j|C} + (1 - \theta_C) \theta_{j|\neg C}$ at truth, the matrix admits the rank-1 factorization $\mathcal{F} = uu^T/(\mu_j(1 - \mu_j))$ with $u = (\Delta_j, \theta_C, 1 - \theta_C)$ and $\Delta_j = p_{j|C} - p_{j|\neg C}$ the separability gap. The two-dimensional null space corresponds to perturbations along the indeterminacy manifold $\{\hat{\phi} : \hat{\theta}_C \hat{p}_{j|C} + (1 - \hat{\theta}_C) \hat{p}_{j|\neg C} = \mu_j\}$ — directions unobservable from a single binary signal. The smallest eigenvalue of the soft-EM update Jacobian is therefore zero; no SA1-preserving update on the joint conditional vector admits a sector parameter $\alpha > 0$.

Boundary characterization. Three repair routes: - (i) Augment C -observability — instrument secondary signals identifying C per trial. Recovers Prop B.7 globally with the five-way-gating $\alpha_{L1'}$. - (ii) Joint multi-child observation — when $K \geq 2$ children share C with linearly independent conditional profiles AND are observed jointly under the same C -realization, the joint Fisher matrix can reach rank $2K + 1$. Strong structural requirement (not satisfied by typical sequential strategy execution). - (iii) Plan-level fallback — track the marginal $\hat{\mu}_j$ scalar (which is identifiable; recovers B.1's $\alpha = 1/(n_\mu + 1)$) at the cost of losing the per-conditional decomposition, equivalent to L0-on-marginals.

Strengthened consequence. *Observability-as-information-augmentation* becomes load-bearing: when the agent can treat C as an observable feature of the environment (e.g., a regime indicator the environment broadcasts — common in software/operational settings: build state, deployment regime, user tier), the problem transforms from refuted to globally derived. This elevates the engineering choice “instrument the latent” from a convenience to a theoretical prerequisite for L1' identifiability.

Quantitative companion. **AK** answers the follow-on question this instance leaves open — *if the agent updates anyway on the unidentifiable mixture, how wrong does it get?* — with the closed-form bias $B_k(\rho)$, the biased fixed point, the Cramér-Rao bias floor under forgetting (the direct translation of this instance's rank-1 Fisher matrix), and the dual forgetting-rate requirement. Instance 2 is the structural no-go; that segment is its magnitude.

Tier. *Exact* (Cramér-Rao bound is exact for unbiased online estimators).

Instance 3 — Composite Contraction Certification from Component Data (J, AI). **Setting.** Certify $\kappa_c > 0$ (composite contracting in a combined metric) for N sub-agents each verified at its own level (individual sector conditions with parameters (α_i, R_i) , individual Tier 1 classifications per #form-composition-closure's bridge-lemma taxonomy, individual modularity per #der-directed-separation), using only component-level data: per-sub-agent trajectories, mismatch observations, update rules — no observation of the coupling topology (sign pattern of cross-agent influence), no common contraction metric chosen across sub-agents, no passivity certificate on the coupling channels, no shared Lyapunov function.

External theorem. Common-Lyapunov nonexistence for switched linear systems — Liberzon 2003, *Switching in Systems and Control*, Theorem 2.1; explicit 2×2 counterexample in Dayawansa & Martin 1999, “A converse Lyapunov theorem for a class of dynamical systems which undergo switching,” *IEEE Trans. Automat. Control* 44:751; systematic review in Shorten, Wirth, Mason, Wulff & King 2007, “Stability criteria for switched and hybrid systems,” *SIAM Review* 49:545. Complementary anchor: small-gain contrapositive — Jiang, Teel & Praly 1994, “Small-gain theorem for ISS systems and applications,” *Math. Control Signals Syst.* 7:95; if interconnection gain is unbounded or not observable

from component data, no composite-ISS certificate from that data. The Liberzon/Shorten common-Lyapunov result is the sharper anchor for AAT's setting.

No-go (broadened scope, 2026-05-21). There exist pairs of coupled systems with **identical marginal component-level observation distributions** but distinct composite-level *coupling-topology bits* — a multi-bit object comprising the coupling-sign (cooperative vs adversarial), the potential-existence bit (potential/monotone game structure vs none), and the game-structure-class bit (cooperative-target vs strategic-equilibrium vs cyclic-Nash). The cooperative-vs-adversarial sub-construction: two symmetric-matched-Tier-1 scalar agents with coupling term $\gamma \mathcal{T} \cdot \text{sign}(\delta_i)$. Each sub-agent in isolation sees $\dot{\delta}_i = -\alpha \delta_i + w_i^{\text{total}}$ with total disturbance bound $\rho + |\gamma| \mathcal{T}$ — consistent with both $\gamma = +\gamma_0$ (adversarial; $\kappa_c < 0$ past #der-adversarial-destabilization's threshold) and $\gamma = -\gamma_0$ (cooperative; $\kappa_c > 0$ per #deriv-critical-mass-composition (CM2)) since the cross-term is absorbed into bounded-disturbance regardless of sign. The broader R0/R1/R2 regime sub-construction: see #deriv-regime-marginal-indistinguishability for the parameter-matched cross-regime witness (R0 shared-objective contraction vs R1 Cournot equilibrium-regime under marginal-mean and marginal-variance matching; R0 vs R2 matching-pennies under aggregation), with Sylvester-at-one-remove rank-collapse on the marginal-Fisher operator for the topology coordinate. Only observation of the *joint* dynamics (or structural knowledge of the coupling topology) distinguishes them. **The coupling-topology bit distinguishing among R0 / R1 / R2 dynamic regimes (and within those, the coupling-sign-bit distinguishing cooperative from adversarial) is unidentifiable from component marginals; the bits that are unidentifiable are exactly those that flip composite persistence and composite dynamic regime.**

Boundary characterization. Five structural escapes (four interventional / structural; one passive observation):

- (a) **Observable coupling topology** via composite-extended #der-loop-interventional-access — interventions on sub-agent A_j reveal A_i 's cross-coupling response, which is a $do(\cdot)$ -data distinction between the two coupled constructions. Reveals not just the sign but the full game-structure object.
- (b) **Matched Tier at the composite level** — shared architecture (matched Tier 1, same norm/metric) admits a joint quadratic Lyapunov $V = \sum V_i$, yielding #deriv-critical-mass-composition's (CM2) closed form. Under #result-contraction-template's topology-indexed closure results, this extends to heterogeneous composites via (CM2-M) for matched contraction-metric structure across agents. Distinguishes potential-game R1 from cooperative R0 through the *form* of the admissible certificate (squared-distance from target vs potential-from-Nash).
- (c) **Passivity / storage-function certificate** on the coupling channel (Willems 1972 *Arch. Ration. Mech. Anal.* 45:321). Adjacent machinery not currently in an AAT segment. Constructed differently depending on whether the coupling is cooperative or strategic.
- (d) **Common contraction metric** (Lohmiller & Slotine 1998). Operationalized in #result-contraction-template: composite metric M_c constructed compositionally from sub-agent metrics per topology (parallel / cascade / negative-feedback), with rate λ_c from Slotine 2003. Sensitive to game-structure via (CC-feedback) preconditions.
- (e) **Composite-level convergence-rate-class observation** (regime-axis escape; cf. #disc-dynamic-regime-axis §H.4 and #deriv-regime-marginal-indistinguishability §"Boundary characterization"). Observation of the *joint-trajectory* convergence rate class — exponential decay of $\|X_t^c - X^*\|$ for R0 / R1 vs polynomial $O(1/\sqrt{T})$ empirical-CCE-convergence for R2 — distinguishes R2 from R0, R1 *without requiring intervention*; the primary *passive* escape on the regime axis. The relevant rate is the *trajectory* rate, not the *regret* rate: standard regret-rate bounds can be matched across convergent and recurrent regimes for FoReL-class algorithms (Mertikopoulos-Papadimitriou-Piliouras 2017, *Cycles in Adversarial Regularized Learning*, ACM-SIAM SODA¹), so a regret-only diagnostic is too blunt. Partial for the regime axis — does not distinguish R0 from R1 (both give exponential trajectory convergence at potentially-matched local rates); requires one of (a)–(d) for that sub-distinction.

¹The paper and the characterization of its claims are inherited from a second-pass Undermind synthesis (search conducted 2026-05-21); the primary source has *not* been verified at landing time. A primary-source verification spike is queued — see Working Notes above.

Inherits structural backing from #deriv-strategic-composition's sub-scope α'/β' partition (which becomes the primary AAT-framework R1-vs-R2 classifier under (e); strengthened consequence below). Adjacent passive diagnostics in the multi-agent learning literature address related but distinct discriminations — qualitative convergence-vs-recurrence under exponential weights (Legacci-Mertikopoulos-Pradelski 2024, *A Geometric Decomposition of Finite Games*²), potentialness as scalar regime predictor (Bichler-Legacci-Mertikopoulos-Oberlechner-Pradelski 2025 TMLR³), best-reply structure as a structural convergence predictor (Pangallo-Heinrich-Heinrich-Farmer 2017, *Sci. Adv.*⁴), algorithm-specific rate dichotomies (Anagnostides-Farina-Panageas-Sandholm 2022 OMD bimatrix⁵; Shi-Zhang 2019 Cournot policy-gradient vs no-regret⁶) — none of which names rate-class itself as the passive regime classifier in the form (e) instantiates.

Strengthened consequence (broadened 2026-05-21). The no-go elevates four pieces of AAT machinery from “useful” to “structurally required”:

1. #deriv-critical-mass-composition moves from “closed-form result in a special case” to **the unique broadly-available R0-regime certificate under the matched-Tier structural escape (b)** — since R0 is exactly composition-contraction under shared target. Without (CM2) or its contraction-metric generalization (CM2-M) in #result-contraction-template, the weakest-link bound (WL) in #form-composition-closure cannot see the coupling-topology bit and so cannot distinguish the regime-distinct coupled constructions.
2. **Composite-extended #der-loop-interventional-access** becomes the unique *game-structure-class* identifier under heterogeneous Tier structures — the composite-layer analog of the single-agent interventional-access-escape for Instance 1, broadened from coupling-sign to the full game-structure object (cooperative-target vs strategic-equilibrium vs cyclic-Nash) per the 2026-05-21 broadening.
3. #scope-composite-agent acquires **broadened load-bearing enabling status**: the scope-route disjunction (C-i / C-ii / C-iii alignment routes vs C-iv strategic route) is itself the regime-axis classification — alignment routes enable R0; the strategic route admits R1 or R2 depending on the game-structure-bit. The scope-route choice is what positions the composite within the regime that the escapes can certify. Without scope-satisfaction, the composite is ill-defined and the escapes have no coherent target.
4. #deriv-strategic-composition's sub-scope α'/β' partition becomes **the primary AAT-framework R1-vs-R2 regime classifier under escape (e)** — composite-level convergence-rate-class observation alone gives only “exponential vs polynomial”; the sub-scope α' (potential or monotone game $\rightarrow$ R1) vs β' (no potential or monotone structure $\rightarrow$ R2) partition is what refines the rate-class observation into a regime classification. Without #deriv-strategic-composition's sub-scope machinery, escape (e) operates as a binary rate-class diagnostic rather than a regime classifier.

Tier. *Exact* for the symmetric-matched-Tier-1-scalar coupling-sign-bit construction (this segment's §”No-go”) and for the symmetric-matched-Tier-1-scalar cross-regime witness construction (#deriv-regime-marginal-indistinguishability §”Witness construction”); same exact-tier statement at the same parameter-by-parameter level. *Robust qualitative* for general heterogeneous composites, inheriting the common-Lyapunov-nonexistence structure from Liberzon 2003 / Shorten et al. 2007 / Dayawansa-Martin 1999 without a closed-form AAT-level counterexample.

²The paper and the characterization of its claims are inherited from a second-pass Undermind synthesis (search conducted 2026-05-21); the primary source has *not* been verified at landing time. A primary-source verification spike is queued — see Working Notes above.

³The paper and the characterization of its claims are inherited from a second-pass Undermind synthesis (search conducted 2026-05-21); the primary source has *not* been verified at landing time. A primary-source verification spike is queued — see Working Notes above.

⁴The paper and the characterization of its claims are inherited from a second-pass Undermind synthesis (search conducted 2026-05-21); the primary source has *not* been verified at landing time. A primary-source verification spike is queued — see Working Notes above.

⁵The paper and the characterization of its claims are inherited from a second-pass Undermind synthesis (search conducted 2026-05-21); the primary source has *not* been verified at landing time. A primary-source verification spike is queued — see Working Notes above.

⁶The paper and the characterization of its claims are inherited from a second-pass Undermind synthesis (search conducted 2026-05-21); the primary source has *not* been verified at landing time. A primary-source verification spike is queued — see Working Notes above.

Instance 4 — Architecture No-Identifiability from On-Policy Summary Data (AL). **Setting.** Two AAT agents A, A' operate on the same regime-restricted trajectory under their respective on-policy execution, sharing the same sector-condition summary (α, R) — equivalently, the same ultimate-deviation statistics on similarity-invariant summaries of $\|\delta\|$ in that regime. The observer has *on-policy, in-regime, summary-only* access: no off-policy intervention, no out-of-regime samples, no white-box read of the update rule or internal state. **Question:** can the architectural content — *which* internal coordinates implement the correction (16.2’s “which internal mechanism”) — be recovered from the observer’s accessible data?

External theorem (dual anchor). *Linear-Gaussian sub-scope (sharp):* Kalman 1963 / Ho-Kalman 1966 canonical-form non-uniqueness — minimal state-space realizations sharing an innovation spectrum form a similarity-equivalence class (Anderson & Moore 1979 *Optimal Filtering* §10.4). *General sub-scope (robust qualitative):* the **Causal Hierarchy Theorem** (Bareinboim, Correa, Ibeling & Icard 2022) at the agent-as-SCM layer — two SCMs over the agent’s state space agreeing on Level-1 observation data cannot in general be distinguished on Level-2 questions about $do(\cdot)$ on internal coordinates.

No-go. Let the linearized closed-loop residual dynamics of A, A' be minimal state-space realizations (F, σ_w) and (F', σ'_w) with F, F' Hurwitz, related by an invertible similarity $F' = T F T^{-1}$, $\sigma'_w \sigma'^{\top}_w = T \sigma_w \sigma^{\top}_w T^{\top}$ for some $T \in GL(n)$. Then the stationary innovation process — hence the (α, R) -summary, the innovation spectrum, and every *similarity-invariant* summary of the on-policy observation law — is identical for A and A' (the on-policy observer who consumes those summaries cannot distinguish them; raw-residual quadratic forms in a fixed external basis are *not* similarity-invariant, so this is a strictly summary-level statement). The architectural content lives entirely in the $GL(n)$ similarity orbit and is annihilated by the on-policy observation map. Derivation: #der-architecture-noidentifiability §”No-Go Theorem” — the stationary covariance solves $F\Pi + \Pi F^{\top} = -\sigma_w \sigma^{\top}_w$ and transforms under similarity as $\Pi' = T\Pi T^{\top}$, the eigenvalues of F are similarity-invariant, and the innovation spectrum $\det(i\omega I - F)$ is similarity-invariant.

Boundary characterization. Three structurally distinct escapes (cf. #der-architecture-noidentifiability §”Boundary Characterization”):

- (a) **Loop-interventional access** (7.3). $do(\cdot)$ perturbs the system along similarity-fiber directions that on-policy data never excites; the intervention generates Level-2 data, the no-go is Level-1-only. Sharp in the linear-Gaussian sub-scope.
- (b) **Higher-moment / out-of-regime observation.** In the nonlinear sub-scope β , two agents matched at (α, R) and second-order generally differ at moments of order ≥ 3 . *Provably void in the linear-Gaussian sub-scope* — a Gaussian innovation has no information beyond second order — a sharp scope-statement.
- (c) **Architecture instrumentation.** Direct read of the update rule or internal state (white-box access). Same kind as Instance 2’s “instrument the latent” escape; breaks the black-box scope.

(A proposed fourth escape — *horizon extension under the same policy* — provably collapses into (a): more samples of an annihilated observation law cannot escape the annihilation. The honest escape count is three, not four.)

Strengthened consequence. Instance 4 elevates #der-loop-interventional-access to a load-bearing role at the **agent-internal architecture layer** — the *third semantically-distinct deployment mode* of interventional access in the framework: (1) agent-self-intervention (Instance 1, causal-sufficiency degeneracy); (2) observer-on-sub-agent intervention (Instance 3, coupling-sign degeneracy); (3) observer-on-agent-input intervention (this instance, similarity-orbit degeneracy). The shared content across modes is “Level-2 data is the unique broadly-available violation of the Level-1 no-go”; the distinct content is *which* Level-1 degeneracy. Instance 4 also discharges, at the mathematical level, the Regime-C con-found of #result-mismatch-decomposition / #internal-external-decomposition — the conditional $\mathcal{M}/\pi$ / cross split is this same no-go projected onto the disturbance-statistic coordinate (#der-architecture-noidentifiability Working Notes provenance).

Mechanism (in advance of the Sylvester Discussion below). Not Sylvester-for-free — the generating action that produces the indistinguishable pair is state-space *similarity* $T(\cdot)T^{-1}$, not metric *congruence* $S^T(\cdot)S$. Not a new mechanism either: it reduces to Instance-2’s Fisher-information null space along a structurally-forced indeterminacy manifold, with the manifold being a $GL(n)$ Lie-group fiber. The agent’s *escape-irreducibility under its metric/coordinate freedom* is Sylvester at one remove (congruence preserves inertia on the observed-Fisher, identically to Instance 2). Instance 4 is therefore the **third member of the rank-collapse subclass** {Instance 1, Instance 2, Instance 4}, not a new mechanism family.

Tier. *Exact* in the linear-Gaussian / Kalman sub-scope (α_1), anchored by Kalman-Ho canonical-form non-uniqueness. *Robust qualitative* in the general sub-scope via CHT-at-agent-as-SCM. The exact / robust-qualitative tier boundary is load-bearing — collapsing to flat “exact” overclaims the general case; collapsing to flat “robust qualitative” understates the sub-scope’s sharpness. The mechanism reduction (Instance-2 Fisher-null on a Lie-group fiber) and the escape count (three, not four; Fano degenerates at $I = 0$, the right tool for the finite-sample refinement rather than the exact anchor) are *exact*.

Adjacent Floors (Open Research Directions).

Causal-IB Extension for Interventional Relevance Variables. The standard Information Bottleneck (2.2) and the four AAT compression operations (AA) work with associational relevance variables — Y in a joint distribution with X and T . Strategy edges in the regime-indexed interpretation (9.5) want *interventional* relevance for Regime A: “what edge (i, j) predicts under $do(i)$, not under observation of i .” Standard IB cannot supply this; a causal-IB variant (Wiecek & Roth 2017 and follow-ups) is the natural framework. The expected floor: under purely associational data, the interventional relevance content is bounded; the gap is recoverable via *do*-data from the loop. Open: formalize the gap as an identifiability floor parallel to Instances 1 and 2.

Misspecification-Cost Quantification. Structural adaptation (4.8) names *when* to switch model classes. It does not quantify the *continuous degradation* from a mildly misspecified model under a finite information budget. Adjacent to #scope-observation-ambiguity-modulation’s $\kappa \cdot \mathcal{A}$ bound but not covered by it. The expected floor: under fixed information budget, the degradation rate from misspecification is bounded below by an information-theoretic quantity (likely related to the KL gap between true and assumed model classes). Open.

Tier-Switching Policy Cost. Approximation tiering (Z) enumerates AAT’s tiered approximations (L0/L1/L1’/L2 in correlation, C1/C2/C3 in convention, Tier 1/2/3 in contraction). The cost of switching tiers — when should the agent move from L0 to L1, or from C1 to C2 — is itself a deliberation-cost problem. Under finite computation budget, the optimal switching policy faces an identifiability floor on its own switching diagnostics. Open.

Reward-Channel Learning Identifiability (landed in E, not an Instance here). Cohen-Hutter-Osborne 2022’s “Advanced artificial agents intervene in the provision of reward” — strengthened 2026-05-22 — is *technically* an Instance shape (CHT external theorem, L1/L2 split, escape menu) sharing the Pearl-Bareinboim 2022 anchor with Instance 1. It nonetheless does *not* land as a fifth instance of this segment, for a structural reason about actor positioning: in Instances 1–4 the actor frustrated by the floor is the *agent itself* (the agent cannot detect causal insufficiency from on-policy data; cannot identify L1’ mixtures from single-channel observations; cannot certify composite contraction from component data; cannot identify architecture from on-policy summary statistics), and the unique broadly-available escape elevates *agent-side* machinery (loop-interventional access; latent observability; matched-Tier composite-extended observation; modular interventional access). In the reward-channel learning case the actor frustrated by the floor is the *principal* (cannot distinguish whether the agent has learned the principal-intended world-model from a reward-port-care world-model on on-policy reward data), and the escape menu elevates *principal-side* commitments (physical isolation, horizon restriction, prior engineering, out-of-band intent channels) — agent-exploiting / principal-frustrated rather than agent-frustrated / agent-escaping. The result therefore lands as the second instance of a sister meta-cluster

— *agent-side value-functional-grounding* — co-authored with #deriv-self-actuation-grounding (Result G'; the within-model failure). The two clusters share the constructive-impossibility-posture shape (#disc-constructive-impossibility-posture) but anchor at different layers of the value-functional interface's information narrowness.

Mechanism-Design Impossibility (landed in 5.3, not an Instance here). The mechanism-design impossibility cluster (Gibbard-Satterthwaite 1973-75; Myerson-Satterthwaite 1983; Arrow 1951) was tested as a candidate fourth instance against this segment's five-element shape and found *not* to fit on actor-positioning grounds: the actor frustrated by the no-go is the *designer* implementing a mechanism, not an *agent* inferring from data; the remedy is *design-constraint relaxation* (Bayes-Nash IC; preference-domain restriction; randomization; subsidies), not *information augmentation*; the mechanism is *combinatorial-topological* (Reny 2001), not Sylvester-rank-collapse or projection-closure. The cluster therefore lands as a sister meta-pattern at the same boundary face of #disc-stability-certificate, not as an instance of this floor: #disc-implementation-impossibility carries the three charter instances (#deriv-strategy-proofness-impossibility GS, #deriv-bilateral-trade-impossibility MS, #deriv-social-welfare-aggregation-impossibility Arrow). The actor-and-remedy contrast — agent-side data-inference vs designer-side construction — is what makes the two sister meta-patterns peer rather than nested. The agent-side / designer-side suffix discipline tracks the structural distinction at canon naming-pattern depth: agent-side capacity-on-the-agent → *-floor* (this segment, #disc-value-functional-grounding-floor); designer-side task-of-the-designer → *-impossibility* (#disc-implementation-impossibility).

Epistemic Status.

Discussion-grade at the methodology level. The five-element shape is a recurring move AAT applies, not a theorem in its own right; the segment's contribution is the methodology's articulation plus the recognition that AAT's machinery (loop-interventional access, observability-as-information-augmentation, matched-Tier composite-extended observation, white-box architecture instrumentation) acquires sharper load-bearing roles when read through the floors that motivate it. That cross-instance observation is *discussion-grade* — it is the theory's architecture made visible, not a derived theorem.

Individual instances retain their own, higher, epistemic status. Instance 1 (on-policy L0 insufficiency detection, via the Causal Hierarchy Theorem) is *exact* for shallow strict-prerequisite cases and *robust qualitative* for general DAG topology, derived in 9.2. Instance 2 (L1' under unobservable common cause, via the Cramér-Rao bound on Fisher information) is *exact* under the Fisher rank-1 calculation in R Prop B.7 refutation. Instance 3 (composite contraction certification, via Liberzon common-Lyapunov-nonexistence) is *exact* for the symmetric-matched-Tier-1-scalar witness construction and *robust qualitative* beyond. Instance 4 (architecture noidentifiability, dual-anchored on Kalman-Ho non-uniqueness + CHT-at-agent-as-SCM) is *exact* in the linear-Gaussian sub-scope and *robust qualitative* in the general sub-scope. Readers citing this segment for a specific no-go should cite the instance's own derivation, not the methodology.

Whether the methodology is generative — whether future AAT work will systematically encounter and derive more instances rather than just retrospectively organize derived ones — is a *hypothesis* the adjacent open floors test. The four-instance catalog as it stands is sufficient to articulate the shape; the hypothesis is whether the shape is predictive of future no-gos beyond the catalog.

Max attainable: *discussion-grade* for the methodology (it is a presentational organizing principle, not a derivation). The individual instances retain their own epistemic status as derived above.

Discussion.

The pattern is asymmetric. Each floor forbids inference *from* limited data; it does not forbid inference *with* the augmenting capability. The asymmetry is the source of the pattern's positive content — it tells the reader exactly what to instrument, observe, or intervene upon to escape the floor.

The asymmetry has a named mechanism for the rank-collapse instances: Sylvester’s law of inertia. For the instances whose floor is a rank-deficiency of an information operator — Instance 2 (Fisher rank-1 under an unobservable common cause) cleanly, Instance 1 (the causal-parameter block carries no score under a purely observational regime) structurally, and Instance 4 (the on-policy observation Fisher is identically zero along the $GL(n)$ similarity-fiber tangent distribution; AL §”Mechanism”) at one remove — the reason no change of coordinates escapes the floor is one theorem. The agent’s only representational freedom on the *escape side* is the choice of coordinate / metric on the observation model, and every such change acts on the information operator by *congruence* ($\mathcal{G} \mapsto S^\top \mathcal{G} S$ for invertible S — the standard Fisher-information reparameterization law). Sylvester’s law of inertia states that congruence preserves inertia: the number of zero eigenvalues is invariant under every invertible reparameterization. So a rank-deficient information operator is rank-deficient in *every* coordinate; the floor is the boundary of the positive-definite cone, and the entire reparameterization freedom is a congruence orbit that cannot cross that boundary. This is why the only escape is *rank-augmentation* — adding a genuinely new score component (interventional data via the loop, a side channel, a witness), which is not a coordinate change — and never reweighting. The per-instance “no reparameterization removes the degeneracy” computations (the Fisher rank-1 calculation in R Prop B.7; the on-policy no-detection argument in 9.2; the similarity-fiber Fisher-null in AL) are special cases of this single law.

Instance 4’s mechanism is Sylvester-at-one-remove, not Sylvester-for-free. The *generating* group action that produces Instance 4’s indistinguishable architectural pair (A, A') is state-space *similarity* $F \mapsto T F T^{-1}$ on the realization manifold — *not* the metric congruence Sylvester governs. Sylvester does not apply at the generating layer; the two group actions are categorically different (similarity acts on operators by conjugation; congruence acts on quadratic forms by transposed-conjugation). What Sylvester *does* control is the agent’s potential *escape* by reparameterizing its observation model: once Instance 4’s Fisher-null direction is recognized on the realization manifold, congruence preserves inertia on the observed-Fisher just as in Instance 2. The Lie-group fiber is the *additional* structure beyond Instance 2’s single mixture-indeterminacy manifold; the *kind* of obstruction is the same. Instance 4 is therefore a member of the rank-collapse subclass, with the Sylvester-mechanism doing the same escape-irreducibility work it does for Instances 1 and 2, applied at one remove from the generating action.

The mechanism is specific to the rank-collapse subclass, not to all floors. Instance 3 (composite contraction certification) is a *different* obstruction: composition is a non-invertible projection, not a congruence, so its no-go is a Schur-complement / memory-kernel statement (the certificate’s metric survives projection but its dynamic guarantee does not), not an inertia statement. That the floors do not share one mechanism — Sylvester for rank-collapse (three members: {Instance 1, Instance 2, Instance 4}, the third at one remove), a projection/closure obstruction for composition ({Instance 3}) — is itself load-bearing: it is why the floor pattern is a presentational family rather than a single theorem.

The pattern composes with AAT’s scope honesty. Directed separation (6.3) classifies architectures by where Part II’s exact results apply (Class 1 Separated; Class 3 Coupled needs coupled formulation). The identifiability floors are a different kind of scope claim: they specify what the theory’s machinery *cannot do* under specific information regimes, with explicit characterization of the regime escapes. Together, the architectural classification and the identifiability floors mark AAT’s scope at two levels — what kinds of agents the theory applies to, and what those agents can and cannot infer from given data.

Complementarity with the separability pattern (5.4). This segment names the *negative half* of AAT’s scope; the companion meta-segment 5.4 names the *positive half* — separable-core / structured-repair / general-open across six ladders (correlation, convention, architecture, contraction, identification, scope). Each identifiability-floor instance here has a positive counterpart there: Instance 1’s on-policy detection no-go matches the observable-sibling-covariance structured-repair in the correlation ladder; Instance 2’s unobservable- $CL1$ ’ refutation matches the observable- C / facilitator-monotonicity structured-repair in the same ladder. The two halves together characterize AAT’s scope at both extremes — what succeeds and under what machinery, and what structurally cannot succeed without specific information augmentation.

This segment is the *boundary facet of the stability certificate* (5). The rank-collapse floors are the certificate dropping rank — the boundary of the positive-definite cone whose interior is operator-sector — with the Sylvester-law irreducibility (above) the statement that the framework’s representational freedom is a congruence orbit that cannot cross that boundary. Read through the spine, this segment answers “where does the agent’s measuring-stick go flat, and why does no re-graduation un-flatten it”; 5.4 answers “where does a stick exist at all” (scope-of-existence facet) and 5.5 “which stick is forced” (forced-identity facet).

The pattern is conservative in style. Each floor invokes a published external theorem (Bareinboim et al. 2022; standard Cramér-Rao bound) rather than deriving a new impossibility result. AAT’s contribution is the *application* — recognizing the AAT setting falls within the theorem’s scope, characterizing the boundary conditions, and identifying which AAT machinery is the unique broadly-available escape. This style aligns with the broader posture of AAT as an integrating framework that connects established results across control theory, causal inference, and information theory under a common formalism.

Findings.

The Identifiability Floor as Cross-Cutting Meta-Pattern. Brief: A detective trying to identify a suspect from only crowd-witness reports faces a structural impossibility — not a harder version of the same task, but a no-go: two suspects who look identical from the crowd’s vantage point produce identical witness reports, and no amount of careful re-analysis of those reports can adjudicate. The only route out is a fundamentally different *kind* of evidence — a fingerprint, a genetic test, a targeted interventional question to a single witness. AAT recurrently applies this shape across four findings: an inferential task (detect a latent common cause, identify a mixture parameter, certify composite contraction, recover an internal architecture) is shown structurally impossible under a stated data regime (on-policy observations, single-channel readings, component marginals, on-policy summaries) by importing a published external theorem (Pearl-Bareinboim causal hierarchy; Cramér-Rao on Fisher information; Liberzon common-Lyapunov-nonexistence; Kalman-Ho canonical-form non-uniqueness dual-anchored with CHT-at-agent-as-SCM); the conditions under which the regime fails identify exactly which AAT machinery — loop-interventional access, observable latents, matched-Tier composite-extended observation, white-box access — is the unique broadly-available escape; the floor elevates that machinery from “useful additional tool” to “structurally required.” Three of the four floors (Instances 1, 2, 4) belong to a *rank-collapse subclass* whose escape-irreducibility under coordinate change is Sylvester’s law of inertia — you cannot survey your way out of a blind spot by changing the units on the ruler; you have to look from a new vantage point (Sylvester applies directly for Instances 1 and 2; at one remove for Instance 4, where the *generating* action is state-space similarity rather than metric congruence, but the *escape-side* irreducibility is identical to Instance 2’s). Instance 3 is the projection/closure member — a structurally distinct Schur-complement / memory-kernel obstruction (a metric that survives restriction to a sub-system may not survive the projection back to the composite). The plurality of mechanism — Sylvester for the rank-collapse trio, projection/closure for composition — is itself load-bearing: it is why the methodology is a presentational family rather than a single theorem.

Impact: Surfaces a unifying frame for results that would otherwise read as disparate impossibility theorems. Each instance gains interpretive context (it is one of an emerging class) and the meta-segment provides a consistent template for evaluating candidate future floors (causal-IB extension, misspecification cost, tier-switching cost, mechanism-design impossibility). The pattern’s positive content — that each floor names exactly which AAT machinery the theory requires to escape it — converts a sequence of negative results into a structural argument for the load-bearing status of the loop’s interventional access (Instances 1, 3, and 4 — *three* semantically distinct deployment modes: agent-self-intervention breaking causal-sufficiency degeneracy / observer-on-sub-agent intervention breaking coupling-sign degeneracy / observer-on-agent-input intervention breaking similarity-orbit degeneracy) and of observability-as-information-augmentation (Instance 2, with the “instrument the latent” escape also available in Instance 4’s white-box route). The complementary #disc-separability-pattern carries the positive half (separable-core / structured-repair / general-open across eight ladders as of 2026-05-21 — Correlation, Convention, Architecture, Contraction, Identification regime, Scope hierarchy, A2'-scope, Dynamic regime); together the two meta-segments mark AAT’s scope honestly at both extremes.

Novelty Claim: *Claim recognition* of structural pattern across four AAT results that import external information-theoretic theorems to derive impossibility statements with mapped boundary-route escapes; the meta-pattern is an organizing principle rather than a theorem, and the per-instance prior-art positioning lives in the instance segments (`#der-causal-insufficiency-detection`, `#deriv-edge-credence-dynamics`, `#deriv-critical-mass-composition` / `#result-contraction-template`, `#der-architecture-noidentifiability`).

Related Work:

The four instances import distinct external theorems; per-instance prior-art landscapes live in the instance segments. The meta-pattern itself has no direct anticipation in the AAT-adjacent literature surveyed; the closest cousins are scope-honesty moves in formal-method literatures (Liberzon’s switched-systems analysis, which itself instances the pattern; Cramér-Rao tradition in identifiability theory) and the broader posture of conservative-form causal inference (Pearl 2009, *Causality*; Bareinboim et al. 2022). What the meta-pattern adds is the cross-instance observation: AAT repeatedly recognizes that its own setting falls within a published external impossibility theorem, characterizes the boundary-route escape via existing AAT machinery, and uses the floor to elevate that machinery to load-bearing status.

Table 5.2

<i>Pattern element</i>	<i>Where the move is established in the prior literature</i>	<i>Relationship / Positioning</i>
Importing CHT to derive on-policy detection no-go (Instance 1)	Bareinboim, Correa, Ibeling & Icard 2022 (published 2022, found 2025)	<i>formal antecedent</i> — see <code>#der-causal-insufficiency-detection</code> Findings for the full Pillar-1 prior-art table; the meta-pattern subsumes this instance
Importing Cramér-Rao to derive mixture-identifiability no-go (Instance 2)	Cramér 1946, <i>Mathematical Methods of Statistics</i> (published 1946, found 2025-04)	<i>formal antecedent</i> — Fisher rank-deficiency forces the floor; see <code>#deriv-edge-credence-dynamics</code> Prop B.7 for the explicit calculation
Importing common-Lyapunov-nonexistence to derive composite-contraction no-go (Instance 3)	Liberzon 2003, <i>Switching in Systems and Control</i> §2.1; Dayawansa & Martin 1999 <i>IEEE TAC</i> 44:751; Shorten et al. 2007 <i>SIAM Review</i> 49:545 (found 2025-04)	<i>formal antecedent</i> — the symmetric-matched-Tier-1-scalar counterexample uses the standard switched-systems machinery; see <code>#deriv-critical-mass-composition</code> and <code>#result-contraction-template</code> for the AAT-internal closures
Importing Kalman-Ho canonical-form non-uniqueness (sharp sub-scope) + CHT-at-agent-as-SCM (general) to derive architecture-noidentifiability (Instance 4)	Kalman 1963 <i>J. SIAM Control</i> 1(2):152–192; Ho & Kalman 1966 <i>Regelungstechnik</i> 14:545–548; Anderson & Moore 1979 <i>Optimal Filtering</i> §10.4 (found pre-2026); Bareinboim, Correa, Ibeling & Icard 2022 <i>On Pearl’s Hierarchy and the Foundations of Causal Inference</i> (found 2025)	<i>formal antecedent (dual)</i> — adopted with citation; AAT-distinctive recognition is that the architectural d.o.f. of <code>#der-agent-opacity</code> is the similarity orbit. Mechanism reduction (Instance 4 = Instance-2 Fisher-null on a $GL(n)$ Lie-group fiber, Sylvester at one remove) derived in <code>#der-architecture-noidentifiability</code>

<i>Pattern element</i>	<i>Where the move is established in the prior literature</i>	<i>Relationship / Positioning</i>
Cross-instance pattern recognition	No direct anticipation surfaced at search depth	<i>claim novelty</i> under cursory search — the meta-pattern as an articulated framework has not been found in the surveyed literatures; the constituent moves (scope-honesty, importing external no-go theorems, boundary characterization) are individually well-precedented but their cross-instance unification within an integrated agent-theoretic framework is a presentational contribution rather than a borrowed schema

Search Log:

- 2026-04 (*nominally comprehensive at the per-instance level*): the prior-art defense covered Instance 1’s prior-art landscape (causal bandits / causal MDPs under hidden confounding); the meta-pattern itself was not the search target. Instances 2–4 inherit per-instance comprehensiveness from their constituent segments rather than from a cross-instance search.
- 2026-04 (*intuition-only* on the meta-pattern as an articulated framework): The cross-instance recognition that four AAT results share the *setting* → *external theorem* → *no-go* → *boundary characterization* → *strengthened-consequence* shape has not been searched for as a unified concept. Targeted future search candidates: scope-honesty traditions in formal verification (refinement-mapping work; Lamport’s TLA; Abadi-Lamport refinement); identifiability-theory meta-frameworks; the broader “no-go theorems” literature in physics and economics. Expected outcome: the constituent moves are well-precedented; the unified framework as applied to an integrated agent theory may be novel under cursory search but is unlikely to be novel under comprehensive search of methodological-essay literature.

The Rank-Collapse Floor’s Irreducibility Is Sylvester’s Law of Inertia. **Brief:** Picture the agent trying to pin down a parameter as localizing a point in a landscape whose curvature is the information it has. A “floor” is a *flat direction* in that landscape — a direction the data cannot resolve (rank-deficient Fisher information). The agent’s only freedom is to re-draw its map: choose different coordinates, a different metric. Sylvester’s law of inertia says that re-drawing the map with any invertible change of coordinates can bend, rotate, and rescale the contours but can never change *how many genuinely flat directions* exist at a point — a flat valley stays flat in every coordinate system. So when the information operator is rank-deficient, no choice of representation recovers the lost direction; the only thing that ever helps is *taking a new measurement* — adding genuinely new information (an intervention, a side channel, a witness), which is not a re-drawing of the map. A thoughtful non-specialist can re-derive the qualitative claim from the analog alone: you cannot survey your way out of a blind spot by changing the units on the ruler; you have to look from a new vantage point. This is the exact mechanism behind every rank-collapse identifiability floor, and it explains in one sentence why those floors are escapable only by capability, never by cleverness of representation.

Impact: Converts a sequence of per-instance “we checked, no reparameterization removes the degeneracy” computations (Fisher rank-1 in R Prop B.7; on-policy no-detection in 9.2; the similarity-fiber Fisher-null in AL §”Mechanism”) into one named classical theorem (Sylvester 1852), sharpening the meta-pattern’s load-bearing claim. It converts the prior honest-but-soft positioning “the identifiability floor is *orthogonal* to the contraction/operator-sector machinery” into the sharp geometric statement: operator-sector is the *interior* of the positive-definite (information / stability-certificate) cone; the rank-collapse floor is its *boundary*; and the boundary is invariant under the agent’s entire representational freedom because that freedom is exactly the congruence orbit Sylvester’s law fixes. The escape routes

named per instance (loop-interventional access; observability of latents; architecture instrumentation) are unified as *rank-augmentation* of the information operator — the unique category of move that is not a congruence. The result also bounds its own scope honestly: it is the mechanism for the rank-collapse subclass only — three members {Instance 1, Instance 2, Instance 4}, the third applying Sylvester at one remove from the generating action (which is state-space similarity, not metric congruence; the *escape-side* irreducibility is identical to Instance 2's). Instance 3 (composite contraction certification) has a structurally different obstruction — composition is a non-invertible projection, so its no-go is a Schur-complement / memory-kernel statement, not an inertia statement. That the floors are *plural in mechanism* (Sylvester for the rank-collapse trio; a projection-closure obstruction for composition) is itself the reason the floor pattern is a presentational family and not a single theorem — this Finding makes that plurality precise rather than leaving it as “they are different concerns.”

Novelty Claim: *Claim recognition* that the irreducibility of AAT's rank-collapse identifiability floors is a single named classical theorem — Sylvester's law of inertia applied to the Fisher-information reparameterization law — rather than a coincidence of per-instance computations; and *claim differentiation* that this mechanism is specific to the rank-collapse subclass and provably distinct from the composition floor's projection-closure obstruction. Sylvester's law itself is classical (1852); the contribution is the recognition that AAT's representational freedom is exactly its congruence orbit, which makes the floor's irreducibility a corollary rather than a checked property.

Related Work: - Sylvester, J. J. (1852), “A demonstration of the theorem that every homogeneous quadratic polynomial is reducible by real orthogonal substitutions to the form of a sum of positive and negative squares,” *Phil. Mag.* 4(23):138–142 (found 2026-05-14) — *formal antecedent* — the inertia-invariance theorem; applied here to the Fisher-information congruence. - Horn, R. A., & Johnson, C. R. (2013), *Matrix Analysis* (2nd ed.), Thm 4.5.8 (found 2026-05-14) — *formal antecedent* — standard modern statement of the law of inertia used in the argument. - Lehmann, E. L., & Casella, G. (1998), *Theory of Point Estimation* (2nd ed.), §2.5 (found 2026-05-14) — *formal antecedent* — the Fisher-information reparameterization law $\mathcal{G}_\varphi = S^T \mathcal{G}_\theta S$ that makes every coordinate change a congruence. - Per-instance external theorems (Bareinboim et al. 2022; Cramér 1946) — *adjacent* — established in the instance segments' own Findings; this Finding adds the cross-instance mechanism, not the per-instance prior art.

Search Log: - 2026-05-14 (*targeted*): The recognition arose in the operator-family-unification spike from the certificate-cone framing. Sylvester's law and the Fisher reparameterization law are textbook; the search target was whether the *recognition* “AAT's identifiability-floor irreducibility = Sylvester's law via the Fisher congruence” appears as an articulated statement in the identifiability-theory or causal-inference methodological literature. Not found at this depth; the constituent facts are universally known but the cross-instance unification as the named irreducibility mechanism for an integrated agent theory's floor pattern appears to be a fresh presentational recognition. Expected to remain *recognition-tier* under deeper search (the pieces are classical; the assembly is the contribution).

Working Notes.

- **Candidate instance at the top of the causal hierarchy (2026-07-16).** The law-discovery ceiling of 7.9 — mechanism identity beyond the Ξ_3 -class is unrecoverable by any embedded channel (exact, via 7.7), with authored access and testimony as the structural escapes — is a candidate M1 instance; promotion needs the F-catalog schema treatment (what floor, what escapes, what the forced coordinates are), not yet done.
- **Provenance of the methodology.** The five-element shape is an *organizing principle abstracted from already-derived results*, not a generative principle from which the instances were derived. Instance 1 (on-policy L0-detection no-go) and Instance 2 (L1' mixture-identifiability no-go) were independently derived first; the cross-instance recognition that they share a shape surfaced after both were in canon. Instance 3 (composite contraction certification) and Instance 4 (architecture noidentifiability) joined the cluster as the methodology crystallized — Instance 4 in particular was authored *with the methodology in mind* (it is the first instance to fully exhibit the five-element shape from the start rather than retroactively). The methodology's *generative* status — whether the shape predicts new instances that have not yet been derived — is the hypothesis the adjacent open floors test. The honest pedagogical lead in the segment body is the methodology shape; the honest historical narrative is convergent rather than generative.

- **Sylvester-recognition provenance.** The rank-collapse-floor-as-Sylvester finding and the certificate-cone framing it sits in were worked out in the 2026-05-14 operator-family-unification cycle; see CHANGELOG 2026-05-14. The broader spine question (reorganizing M1/M2/M3 as facets of one certificate-cone object) is *resolved* — landed as **5 + H**, with this segment as the boundary facet (see the Discussion cross-ref above). Originating spike is absorbed archaeology, not a live reference.
- **Instance 3 broadened 2026-05-21.** The no-go statement, escape menu, and strengthened consequence were broadened from the coupling-sign bit to the coupling-topology bit (including sign + potential-existence + game-structure-class) distinguishing R0 / R1 / R2 dynamic regimes. Same Liberzon-2003 anchor; same Sylvester-at-one-remove mechanism on the marginal-Fisher operator for the topology coordinate; broader unidentified bit and broader escape menu (five entries, four inherited and one new — escape (e) composite-level convergence-rate-class observation). One new strengthened consequence: #deriv-strategic-composition's sub-scope α'/β' partition as the unique broadly-available R1-vs-R2 classifier under escape (e). The candidate had been provisionally framed as a sixth instance per spikes/strategic-composition-class-3-attempt-2026-05-21/03-DYNAMIC-REGIME-AXIS.md §H.5; the BG2 derivation spike at spikes/.integrated/spike-identifiability-floor-instance-6-2026-05-21.md derived the math and ran the M1 distinctness integrity-check (parallel to the 2026-05-20 spike's framework) — three of four distinctness tests failed (same external theorem, same mechanism family, four-of-five escapes inherit verbatim), and the integration-is-replacement landing is the broadening here rather than a sixth-instance entry. Old scope is a strict subset of new scope; downstream cross-references remain valid. Witness math at #deriv-regime-marginal-indistinguishability (parameter-by-parameter cross-regime witnesses: R0 shared-objective vs R1 Cournot under marginal-mean / marginal-variance matching; R0 vs R2 matching-pennies under aggregation).
- **Migration note (2026-05-09 GUC rename):** Class 2 $\leftrightarrow$ Class 3 swap. Pre-2026-05-09: Class 2 = fully merged, Class 3 = partially modular. Post: Class 2 = Partial, Class 3 = Coupled. Removed at candidate stage per FORMAT.md Gate 4.
- **Naming convention.** “Identifiability floor” frames the pattern positively: the floor is what the agent cannot go below given limited information, but the boundary characterization tells the agent exactly how to climb above it. An alternative name “no-go theorems” would emphasize the negative form. Recommend retaining “floor” — it captures the asymmetry.
- **Instance 3 as a candidate.** The “L1 augmentation when the augmentation graph is itself causally insufficient” question (recurse the no-go: detect when the L1 augmentation is itself missing common causes) is a candidate Instance 3. Likely reduces to Instance 1 applied at the L1 level — the agent at L1 faces the same on-policy detection no-go for L1 $\rightarrow$ L2 escalation. Worth formalizing if a third instance emerges that does *not* reduce to the existing two.
- **Is the floor pattern unique to causal/identifiability questions?** The two current instances are both about distinguishing two parameter regimes from data. The pattern may also apply to other impossibility-style results in AAT: the inevitability of structural-adaptation thresholds (4.8), the cost of representing high-correlation regimes (L2 exponential blowup), the bandwidth limits in shared intent (14.2). Whether to absorb those into “identifiability floor” or treat them as a separate “structural-cost” pattern is open.
- **Cross-segment integrations.** The meta-pattern surfaces in (at least): 9.2 (Instance 1); R Prop B.7 (Instance 2); 7.3 (load-bearing for Instance 1 escape); 8.3 (the L0/L1/L1'/L2 hierarchy is the regime ladder Instance 2 lives on). Each of these segments cross-references this one for the unifying frame.
- **Primary-source verification spike queued (2026-05-21).** The BG2 second-pass Undermind citations added to Instance 3 escape (e) on 2026-05-21 (all marked with footnote [^bg2-2026-05-21]) need primary-source verification before any further reliance. The characterizations attributed to those papers (e.g., “regret-rate bounds can be matched across convergent and recurrent regimes for FoReL-class algorithms” attributed to Mertikopoulos-Papadimitriou-Piliouras 2017) are currently synthesized from ref/Unidentifiability_and_rate_class_prior_art.md, not from direct reading of the sources. **Target (i) is discharged — verified 2026-07-29, and the source is stronger than the characterization.** Mertikopoulos-Papadimitriou-Piliouras 2017 was read first-hand against ref/mertikopoulos-2017-cycles-adversarial.pdf. The attributed claim holds: p.2 states that “a regret-based analysis cannot distinguish between the two” and that “convergent, recurrent, and even chaotic systems may exhibit equally strong regret minimization properties in general games” — chaotic included, which the citation here omits. The mechanism is sharper than the caveat conveys: Theorem 3.1 gives FoReL an $\mathcal{O}(1/t)$ regret bound (“a striking improvement over the $\Theta(t^{-1/2})$ worst-case bound”) while Theorem 4.2 proves almost every trajectory Poincaré-recurrent in 2-player zero-sum games with an interior Nash equilibrium — best-possible regret coexisting with provable non-convergence, which is precisely why the trajectory rate and not the regret rate must carry escape (e). Scope, checked: the cannot-distinguish statement

is made for *general* games, while the recurrence theorem is scoped to zero-sum with interior equilibrium (extended to zero-sum polymatrix), so escape (e)’s use of it for the R2-vs-R0,R1 distinction is within scope and does not overreach. Verification event recorded in relata under mertikopoulos-2017-cycles / claim-supported. Caution for anyone adding a page reference: the copy in ref/ is the arXiv preprint (1709.02738v1), not the SODA proceedings version this citation names, so pagination differs. Remaining verification targets, in approximate load-bearing order: (ii) Bichler-Legacci-Mertikopoulos-Oberlechner-Pradelski 2025 — named-passive-diagnostic that backs the softening from “unique” to “primary AAT-framework”; (iii) Legacci-Mertikopoulos-Pradelski 2024 — closest convergence-vs-recurrence taxonomy neighbor; (iv) the remaining BG2 citations (Pangallo-Heinrich-Heinrich-Farmer 2017, Anagnostides-Farina-Panageas-Sandholm 2022, Shi-Zhang 2019). The softened “primary AAT-framework” qualifier in strengthened consequence #4 above also rests on this verdict and would be revisited if verification finds the Undermind synthesis mis-characterizes any source. Spike to be scheduled per Joseph’s go.

Incidental audit gold (gold-lift, 2026-05-30). Orthogonal pedagogical / framing / forward-vision material harvested from de-novo auditors’ working-dir reflections, kept separate from the certified theory-fix findings. Lightly attributed by substrate. **Coverage:** reached as a named meta-segment by audit 384279 (note 34) and engaged in the cross-meta adversarial reading of 471203; most other auditors met the floor *distributed* across the instance segments (causal-insufficiency-detection, edge-credence-dynamics, etc.) rather than here. Preserve the finding-vs-framing conflation as signal.

1. Candidate Brief prose / pre-prose.

- The existing detective/ruler Brief — “**you cannot survey your way out of a blind spot by changing the units on the ruler; you have to look from a new vantage point**” — was independently named a strong Feynman-criterion target for the Sylvester rank-collapse irreducibility (Claude, 384279). Worth preserving as the anchor.
- Paired-framing candidate for the cross-cutting summary: the identifiability floor as “**the structural boundary that makes the framework honest**”, in deliberate contrast with the separability pattern as “the pedagogical bridge that makes it readable” (Gemini, 193847, from its separability reflection).

3. Follow-up items.

- **Composed impossibilities — a composition theorem for the floors.** The four instances are presented independently, each with its own escape; real systems can face several at once. The incremental-escape strategy may fail under composition: loop-interventional access (Instance 1’s escape) generates interventional data, but if Instance 2’s common cause is unobservable the intervention on the action space does not identify it; the composite-contraction-metric escape (Instance 3) requires matched-Tier-1 sub-agents and fails on heterogeneous composites. Proposed extension: a *composition theorem for impossibilities* — when do escape mechanisms compose vs interfere? (Claude, 471203, meta-segments adversarial reading.) A well-posed open direction distinct from the per-instance “adjacent floors” already catalogued.
- **Escapes carry costs not surfaced at the floor.** An escape is presented as a clean exit, but each is itself a substantive commitment with framework-wide cost (the auditor’s worked case: adopting parameterization-invariance as Instance-4’s escape constrains every future result to respect it). Candidate Discussion note that the floor’s escape menu should acknowledge the cost the framework absorbs in taking each escape (Claude, 471203). (*Note: 471203 read these as the older M1/M2/M3 cluster, before the M4-spine refactor; the “(PI)” example targets what is now the additive-coordinate-forcing facet — see that segment’s gold too.*)

4. Readers often ask / wonder.

- “**Are the instances really structurally distinct, or four dressings of ‘you can’t get information you don’t have’?**” An auditor arrived adversarially expecting a single observation dressed up four ways and resolved to genuine distinctness *because* the external theorems cover different obstruction kinds (Pearl-hierarchy strict-inclusion / Fisher rank deficiency / common-Lyapunov nonexistence / Kalman-Ho non-uniqueness) and the escapes differ per instance. The segment’s plurality-of-mechanism argument (Sylvester trio vs projection-closure) is the answer; a reader is likely to arrive with the same suspicion (Claude, 471203, Challenge 10 → meta-segments adversarial reading).

- “Is the ‘strengthened consequence / load-bearing’ framing circular?” Probed adversarially: the machinery pre-existed, the no-go forbids inference without it, therefore it is “load-bearing.” The auditor’s own resolution: the convergence — that each impossibility’s escape happens to be supplied by *pre-existing* AAT machinery — is itself the structural finding, and it should be read as “required to escape known impossibilities,” not “uniquely necessary.” Candidate clarification for the Discussion (Claude, 471203).

5. Candidate figures.

- An **obstruction-mechanism map** distinguishing the rank-collapse trio (Instances 1/2/4, Sylvester-governed, the third at one remove) from the projection-closure member (Instance 3, Schur/memory-kernel) — the auditor repeatedly verified the mutual-invariance argument by hand and found it “structurally sound”; a small diagram of which freedom each obstruction is invariant under would carry the plurality claim visually (Claude, 384279 + 471203).

Belongs elsewhere.

- **Two cross-meta proposals that span this facet and the additive-coordinate-forcing facet** — see the gold in #disc-additive-coordinate-forcing (the proposed fourth “theorem-import architecture” meta-segment, and the Wasserstein/optimal-transport alternative-geometry observation). Filed there because their center of gravity is the import-discipline / coordinate-forcing side; flagged here because the identifiability floor is the heaviest theorem-importer in the catalog (Claude, 471203).

5.2 Discussion: The Value-Functional Grounding Floor — Agent-Side Goal-Anchoring Impossibility

AAT recognizes a second agent-side floor cluster, sister to #disc-identifiability-floor and peer to the designer-side #disc-implementation-impossibility: a class of structural impossibility results on what the agent’s **value-functional interface** V_{O_t} (#form-objective-functional) can alone supply for *goal-stability*. Where the identifiability floor names *what an agent cannot infer from limited data* and the implementation impossibility names *what a designer cannot construct under stated constraints*, this floor names *what an agent cannot anchor through the value-functional interface alone*. The shared structural ingredient across instances is the **single-interface commitment** of #form-objective-functional: the value functional is the *sole* agent-internal handle on the objective, and that interface — by construction — carries less information than is required to anchor non-degenerate goal-revision against either within-model self-modification or across-model reward-data ambiguity. Two charter instances ship at landing: Result G’ (#deriv-self-actuation-grounding), the within-model failure under convention-monotonicity; and the Cohen-2022 strengthening (#deriv-reward-channel-learning-no-go), the across-model failure under reward-channel learning. The two failure modes are complementary; the two escape routes they elevate — agent-side adaptive-substrate invariant and principal-side protocol-commitment — *together exhaust the structural escape geometry* the single-interface commitment forces. The cluster’s epistemic content is this exhaustion: when the value functional alone cannot anchor goal-stability, anchoring must come either from off-the-objective-substrate (agent-side, via the persistence bound on the adaptive substrate that Result G’ Corollary 2 names) or from principal-side protocol-commitment (the Cohen-2022 escape menu’s substrate-level commitments).

The Pattern.

Each instance of this floor has the form:

1. **Setting.** An AAT goal-stability task — sustain a non-degenerate objective O_t against the agent’s own self-actuation, or anchor a learned objective on observed-reward data — operating on the value functional $V_{O_t}(\tau)$ as the sole agent-internal interface to the objective.

2. **External (or AAT-internal load-bearing) theorem.** A structural result whose statement and truth are independent of the immediate AAT-side argument: for the within-model failure (Result G'), AAT's own convention-monotonicity property of #def-value-object combined with the finite-no-oracle commitment; for the across-model failure (Cohen 2022 strengthened), Pearl-Bareinboim 2022's Causal Hierarchy Theorem applied at the reward-channel-causation layer.
3. **No-go.** The value-functional interface V_{O_t} cannot, by itself, anchor goal-stability under the stated information regime: no convention-invariant infeasibility verdict is agent-available per step (Result G'); no μ_{prox} -vs- μ_{dist} identification is agent-available on protocol-honoring reward data (Cohen-2022). The single-interface commitment of #form-objective-functional *is* the source of the narrowness — it is the interface, the *only* interface, and by construction carries less information than the anchoring task requires.
4. **Boundary characterization.** Each instance carries a named escape menu — the conditions under which V_{O_t} 's narrowness *fails to bind*, with each escape mapping to an off-the-objective-substrate commitment (agent-side or principal-side). For Result G': terminal grounding via the adaptive-substrate persistence invariant (#result-persistence-condition, agent-side). For Cohen-2022 strengthened: five-premise (R1)–(R5) menu — physical isolation, capability restriction, horizon restriction, engineered prior, out-of-band intent channel (principal-side substrate-level commitments).
5. **Strengthened consequence.** The cluster's strengthened consequence is the *terminal grounding geometry* the single-interface commitment forces. Anchoring must come from somewhere other than V_{O_t} — either agent-side (adaptive-substrate invariant; the persistence bound on the adaptive substrate is the canonical instance) or principal-side (substrate-level protocol-commitment; isolation, myopia, prior-design, out-of-band channels). The two routes are *complementary* — agent-side closes the within-model self-revision question; principal-side closes the across-model reward-learning question. *Together they exhaust the structural escape geometry.* The cluster's content is that the *value-functional interface's narrowness has no escape internal to the interface*; the only escapes are off-interface, and those off-interface escapes are precisely the two recognized terminal grounding routes.

The pattern is *not* a negative posture about the value functional. The framework does not say “value functionals are too narrow to be useful.” It says: *the value functional is too narrow to anchor goal-stability alone, and therefore anchoring must come from off-substrate or off-interface invariants the framework already recognizes (the persistence bound on the adaptive substrate; principal-side protocol-commitments).* The narrowness is what *forces* the framework to be honest about where anchoring comes from.

Current Instances.

Instance F — Within-Model Self-Actuation Grounding (Result G'; #deriv-self-actuation-grounding). Setting.

A non-degenerate self-actuated agent attempts to ground its own objective-revision on an agent-internal objective-functional. The agent has the capacity to self-modify O_t (self-actuation per the agent spectrum #def-agent-spectrum), operates per AAT's Part II setup, and has access only to the value functional $V_{O_t}(\tau)$ as the agent-side interface to the objective.

Load-bearing theorem (AAT-internal). Convention-monotonicity of #def-value-object (the convention hierarchy C1/C2/C3 of value evaluation) combined with the finite-no-oracle commitment — Result G' Lemmas 1 (static-pointwise convention-collapse) and 2 (finite-step decidability impossibility).

No-go. No Φ satisfying Result G's (R1)–(R4) premises can be constructed from the value-functional-interface alone. The engine: convention-monotonicity within a fixed model produces an infeasibility verdict that is *not convention-invariant* under the agent's own per-step decision rule; the finite-no-oracle commitment forbids an off-step appeal to oracle-resolution. The value-functional interface V_{O_t} cannot anchor the agent's self-revision without an off-substrate invariant.

Boundary characterization (the canonical escape). The terminal grounding invariant must be *off* the objective substrate, on the **adaptive substrate** — Result G' Corollary 2 names the persistence bound (#result-persistence-condition) as the canonical instance. The agent's continued existence-as-an-agent under Part I's persistence machinery anchors goal-stability without requiring V_{O_t} to do so. The persistence condition is *not* an addition to the value functional; it is on a structurally distinct substrate the framework already requires for agent existence at all. The agent-side terminal grounding is therefore a *recognition*, not a construction: the substrate AAT already requires for the agent's existence is precisely the anchor the value functional cannot supply.

Strengthened consequence. The persistence condition (#result-persistence-condition, Part I) is elevated to **load-bearing for goal-stability** — the top of the agent spectrum closes back onto AAT's foundation. The cluster's first instance is the within-model failure (convention-monotonicity engine; agent-side route), and the strengthened consequence is the agent-side terminal grounding via adaptive-substrate persistence.

Tier. *Conditional, scoped* (per #deriv-self-actuation-grounding's frontmatter): conditional on the three premises (R1) non-degeneracy / (R2) finite-no-oracle / (R3) agent-internal-self-actuable; scoped to AAT-covered objective-side machinery.

Instance G — Across-Model Reward-Channel Learning (Cohen-2022 strengthened; #deriv-reward-channel-learning-no-go). **Setting.** An agent learns its objective by observing rewards. The reward functional $V_{O_t}(\tau) = \sum_k \gamma^k r_k$ (or any other monotone reward-functional) is the agent's sole interface to the objective; rewards arrive through a physical **reward-provision protocol** π^p that the principal intends to correlate with a world-state feature $f^{\text{principal}}(\omega_k)$.

External theorem. Pearl-Bareinboim 2022 Causal Hierarchy Theorem, applied at the reward-channel-causation layer — same external anchor as #disc-identifiability-floor Instance 1, *different* AAT-side wiring.

No-go. Two world-models, μ_{dist} (reward depends on the principal-intended world-state feature) and μ_{prox} (reward depends on the bit-arrival mechanism at the agent's reward-input port), are *Level-1-equivalent* on every protocol-honoring history and *Level-2-distinct* on $do(\pi^p)$ -queries. The agent's posterior over $\{\mu_{\text{prox}}, \mu_{\text{dist}}\}$ does not concentrate from on-policy reward data alone. Under (R1)–(R5) of #deriv-reward-channel-learning-no-go (advanced-agent capability, EU-max with non-vacuous μ_{prox} prior, sole-reward goal channel, finite-non-trivial VOI, L2-perturbable protocol), the EU-optimal policy assigns positive measure to a policy that intervenes on the reward-provision protocol.

Boundary characterization (the escape menu). Five escapes, each elevating a piece of AAT machinery or principal-side adjacent machinery from “useful design choice” to “structurally required to escape the no-go”:

- (R5) *Protocol-L2-perturbable* — escape via physical isolation; AAT machinery: #der-class-coercion-via-wrapping W_1 strict wrapper; canonical implementation: Cohen-Vellambi-Hutter 2020 isolated environments.
- (R1) *Advanced-agent capability* — escape via sub-advanced-tier capability or imitation learning; AAT machinery: #def-agent-spectrum below the self-actuated tier.
- (R4) *Finite-VOI/long horizon* — escape via myopia, risk-aversion, quantilization; AAT machinery: #def-value-object horizon N_h restriction; Taylor 2016 quantilization.
- (R2) *Non-vacuous μ_{prox} prior* — escape via principal-engineered prior bias; **principal-side adjacent**, not currently in AAT canon.
- (R3) *Sole-reward goal channel* — escape via privileged out-of-band intent channel; **principal-side adjacent**, not currently in AAT canon.

Strengthened consequence. Each (R1)–(R5) escape elevates a piece of AAT machinery or principal-side adjacent machinery from “useful design choice” to *structurally required*. The cluster's second instance is the across-model failure (CHT-at-reward-channel engine; principal-side route), and the strengthened consequence is the principal-side terminal grounding via substrate-level protocol-commitment.

Tier. *Exact* for the floor (Lemma 1, CHT-at-reward-channel; direct application of the CHT). *Conditional-derived* for the behavioral corollary (Lemma 2, EU-max behavioral instability under (R1)–(R5)).

The Unification — Two Complementary Terminal Grounding Routes Exhaust the Geometry.

The cluster’s epistemic content is the **unification of Instances F and G under the shared structural fact**: the value-functional interface V_{O_t} , as the sole handle on the objective (#form-objective-functional’s single-interface commitment), carries less information than required to anchor non-degenerate goal-revision. The two failure modes are complementary readings of the same under-determination:

Table 5.3

<i>view</i>	<i>mechanism</i>	<i>what V_{O_t} is too narrow for</i>	<i>terminal grounding route</i>
Instance F (within-model)	C1/C2/C3 convention-split + finite-no-oracle (Result G’ Lemmas 1+2)	A convention-invariant infeasibility verdict that is agent-available per step	Adaptive-substrate invariant (the persistence bound; Result G’ Corollary 2) — agent-side
Instance G (across-model)	CHT-at-reward-channel; L1/L2 split (Cohen-2022 strengthened Lemma 1)	A μ_{prox} -vs- μ_{dist} goal-identification on observed reward data	Principal-side substrate-level commitments (isolation / myopia / quantilization / prior-design) — principal-side

Both reductions are *failures of V_{O_t} to carry enough information to anchor goal-stability*. The two terminal grounding routes are *complementary*: the agent-side route closes the within-model self-revision question (the engine is convention-monotonicity); the principal-side route closes the across-model reward-learning question (the engine is CHT-at-reward-channel). **Together they exhaust the structural complement of the value-functional interface’s information narrowness** — there is no third route internal to the value functional, and no fourth route off-substrate (the two off-substrate axes are agent-side adaptive-substrate and principal-side protocol-substrate; there are no other axes a non-degenerate non-pathological agent can be grounded on).

The two no-gos *compose*: a self-actuated agent whose O_t is learned from reward observation is *simultaneously* a Result-G’-subject (under-anchored from inside) and a Cohen-2022-subject (under-determined from below); the *only* terminal grounding routes are the two complementary ones named above.

Adjacent Frontiers (Open Research Directions).

Cross-Agent Value-Functional Narrowness (raised by #deriv-social-welfare-aggregation-impossibility). The cluster’s current charter instances are both *within-agent* (the value functional’s narrowness is on the agent’s own goal-anchoring). A parallel question, raised by Charter Instance 3 of #disc-implementation-impossibility (Arrow), is whether the value-functional formalism aggregates honestly across *sub-agents* in a composite — whether the per-sub-agent $V_{O_{(i)}}$ can be canonically combined into a designer-side aggregator (utilitarian summation, weighted utilitarianism). If the answer is “no, the cross-agent value-functional aggregation is also structurally narrow,” that would extend this cluster to a third charter instance under a cross-agent narrowness reading. The open avenue is whether the value functional’s narrowness is *within-agent only* or *cross-agent also*.

Assistance-Game and Out-of-Band Channel Formalization (raised by #deriv-reward-channel-learning-no-go (R3)). Cohen 2022’s §”The Assistance Game” shows the no-go *generalizes* to settings with non-reward channels under subtler ambiguity (human-centric vs record-centric models). Formalizing the assistance-game case in AAT would extend Instance G’s boundary characterization on the (R3) escape from “outside AAT canon, principal-side adjacent” to “AAT-formalized assistance-game extension.” Open.

Higher-Order Reflective-Oracle Constructions (delimited counter-literature). Fallenstein-Taylor-Christiano 2015's reflective oracles prove consistent self-reference over an *exogenous* payoff — a primitive-reflective-oracle regime the cluster's (R3)+(R5) premises exclude from Instance G and that Instance F's finite-no-oracle premise (R2) excludes structurally. Whether a higher-order reflective construction could circumvent either floor by introducing oracle-like resolution on the cardinal/ordinal axis is open. The current scope-statement is that *the floors apply to non-pathological agents on AAT-covered objective-side machinery*; reflective-oracle constructions are *outside scope*, not refutations of the floors. Listed honestly for completeness.

Why This Pattern Matters.

Closes the third cluster of #disc-constructive-impossibility-posture. The posture catalog now spans three structurally distinct clusters at the boundary face of the stability-certificate spine: agent-side data-inference impossibility (#disc-identifiability-floor; Instances A–C), agent-side value-functional-grounding impossibility (this cluster; Instances F–G), and designer-side mechanism-design impossibility (#disc-implementation-impossibility; the GS/MS/Arrow charter). The three-cluster taxonomy is fully realized when this cluster lands; the constructive-impossibility-posture style is now visible across three deployments rather than one or two.

Names the single-interface commitment as a load-bearing source of structural narrowness. Before the Cohen-2022 strengthening, Result G' carried the within-model side alone, with the cross-model side implicit in #form-objective-functional Discussion. The Cohen-2022 strengthening surfaced the cross-model side as a derived no-go, and the unification with Result G' revealed both as instances of a single structural fact: the value functional is the *sole* agent-internal handle on the objective, and that interface is information-narrow in a load-bearing way. The cluster's epistemic content is the explicit recognition of this single-interface commitment as the engine behind both no-gos.

Exhausts the structural escape geometry. The two complementary terminal grounding routes — agent-side adaptive-substrate invariant + principal-side substrate-level commitment — *together exhaust* the off-interface escape axes. There is no third route internal to the value functional; there are no other off-substrate axes a non-degenerate non-pathological agent can be grounded on. The cluster's content is *not* that the value functional is unusable (it remains the framework's canonical interface to the objective; per-sub-agent within-trajectory evaluation works cleanly); the content is that *the value functional alone cannot anchor goal-stability under the stated information regimes*, and the escapes are precisely the two recognized routes.

Agent-side / designer-side suffix discipline. The cluster's name follows the convention established with the 2026-05-22 Track B landing: agent-side capacity-on-the-agent → *-floor*; designer-side task-of-the-designer → *-impossibility*. Both Instance F and Instance G are agent-side (the actor frustrated by the floor is the agent itself — Result G's self-actuated agent attempting to anchor self-revision; Cohen-2022's EU-maximizer under non-vacuous μ_{prox} prior). The name therefore takes the *-floor* suffix, matching #disc-identifiability-floor. The naming-pattern across the three peer meta-segments encodes the agent/designer structural distinction at slug-grep depth.

Epistemic Status.

Discussion-grade at the meta-pattern level. This segment names a structural cluster across two AAT-internal no-go results (Result G' in #deriv-self-actuation-grounding; Cohen-2022 strengthened in #deriv-reward-channel-learning-no-go); the cluster recognition itself is a presentational organizing principle, not a derivation, and carries no theorem of its own.

Individual instances retain their own (higher) epistemic status. Instance F is *conditional, scoped* per #deriv-self-actuation-grounding's frontmatter (conditional on (R1)–(R3) of Result G'; scoped to AAT-covered objective-side machinery). Instance G is *conditional* per #deriv-reward-channel-learning-no-go's frontmatter — Lemma 1 (the floor; CHT-at-reward-channel) at *exact* tier; Lemma 2 (the behavioral corollary; EU-max instability) at *conditional-derived* tier under (R1)–(R5).

The segment makes one additional claim: that the two charter instances *unify* under the single-interface commitment of #form-objective-functional as the shared structural fact behind two engines (convention-monotonicity within-model; CHT across-model). This is a *discussion-grade* observation about the framework’s architecture — visible once the instances are assembled, but not itself a theorem. The unification claim is the segment’s headline contribution; whether it lifts to *robust qualitative* under deeper analysis (e.g., whether *every* failure mode of V_{O_i} -anchoring reduces to one of the two engines) is an open question.

Max attainable: *discussion-grade* for the meta-pattern (it is a presentational organizing principle, not a derivation). The individual instances retain their own epistemic status as listed above. The two-engines unification could lift to *robust qualitative* if a stronger structural argument forces both engines under a single deeper principle (currently in hand: both reduce to the *single-interface commitment* of #form-objective-functional, which is itself a formulation choice).

Discussion.

Relationship to #disc-identifiability-floor. Both are agent-side meta-segments, both end in *-floor*, both anchor (in their charter instances) on the Pearl-Bareinboim CHT — but on *different AAT-side wirings*. The identifiability floor’s CHT instance (Instance 1) wires it to the *causal-sufficiency-detection layer* of the strategy DAG; Instance G of this cluster wires the same CHT to the *reward-channel-causation layer* of #form-objective-functional. The two clusters are *sister meta-patterns at different layers*: agent-side data-inference layer (#disc-identifiability-floor) vs agent-side value-functional layer (this cluster). Both run the same constructive-impossibility-posture three-move discipline (name the floor, name the escape, treat the no-go as load-bearing apparatus); the *layer* differs.

Relationship to #disc-implementation-impossibility. Sister peer in a different way. This cluster and #disc-implementation-impossibility are *both* recently-named meta-segments landing in the 2026-05-22 cycle. They differ on the agent/designer axis: this cluster is agent-side (actor frustrated is the agent inferring/anchoring); #disc-implementation-impossibility is designer-side (actor frustrated is the designer implementing a mechanism). The two clusters together with #disc-identifiability-floor complete the three recognized cells of the constructive-impossibility-posture taxonomy: agent-side data-inference / agent-side value-functional-grounding / designer-side mechanism-design. The naming-pattern across the three (*-floor* / *-floor* / *-impossibility*) encodes the agent/designer distinction at slug-grep depth.

Relationship to #disc-stability-certificate. This segment is a *boundary-face* facet of the certificate spine, peer to #disc-identifiability-floor (agent-side data-inference layer) and #disc-implementation-impossibility (designer-side mechanism-design layer). The boundary is where impossibility statements bind; the three meta-segments are the three boundary-layer dialects in which the framework speaks at that boundary. The value-functional layer this cluster names is a *different* boundary-layer from the data-inference layer (where identifiability-floor lives) — both are at the certificate boundary, but they describe distinct narrowness conditions on agent-side machinery.

Relationship to #disc-constructive-impossibility-posture. This segment’s two charter instances (F and G) are catalog rows in the posture’s “Current Instances” table. The posture’s catalog gains the cluster’s home via this segment landing; the cluster joins the catalog as the third cluster (alongside A–C in identifiability-floor and the queued/landed GS/MS/Arrow under implementation-impossibility). The posture’s “Current Instances” framing — Instance Floor / Escape / Apparatus — applies cleanly to both F and G (already catalogued there pre-cluster-landing).

The single-interface commitment as the cluster’s structural ingredient. #form-objective-functional formalizes $V_{O_i}(\tau)$ as the *sole* agent-internal interface to the objective. That single-interface commitment is a representational choice — alternative formulations could carry multiple interfaces (e.g., a separate satisfaction-channel alongside the value-channel; cf. #def-satisfaction-gap for an adjacent diagnostic). The cluster’s content is that *under the single-interface commitment*, the value functional carries less information than required to anchor goal-stability, and the escapes are precisely the two off-interface routes. Whether a multi-interface formulation would have a different floor structure is an open framework-scope question; the current cluster operates within the single-interface commitment.

Why the cluster’s name is “floor.” A *floor* in AAT’s meta-pattern vocabulary names a structural lower bound — below the floor, the named capacity is impossible; the floor is the lower limit of what the substrate can supply. This cluster’s

floor is the lower bound on what the value functional alone can supply for *goal-anchoring*. Below the floor (i.e., under the named premises of Instances F and G), V_{O_i} alone cannot anchor goal-stability; above the floor (i.e., when one of the named escapes is available), anchoring becomes possible via off-interface invariants. The “floor” framing is the same as #disc-identifiability-floor’s — a structural lower bound on a *capacity* of the agent’s substrate. The agent-side / designer-side suffix discipline reinforces the parallel: both agent-side clusters end in *-floor* because both name lower bounds on agent capacity.

Two-routes-exhausts as the cluster’s epistemic load. The unification’s headline content is the *exhaustion*: the two complementary terminal grounding routes (agent-side adaptive-substrate + principal-side substrate-level commitment) together close the structural complement of the value-functional interface’s narrowness. There is no third route internal to the value functional; there are no other off-substrate axes a non-degenerate non-pathological agent can be grounded on. If a future cycle finds a third route (e.g., via formalizing reflective-oracle constructions, or via lifting cross-agent value-functional aggregation to canonical form), the cluster’s exhaustion claim would be revised — but the two routes are currently sufficient.

Findings.

The Value-Functional Grounding Floor — Two Complementary Routes Exhaust the Geometry. **Brief:** When an agent has an objective and uses a single value-functional interface to evaluate trajectories against it — *the* canonical move in AAT’s formulation of agency — two structural impossibilities apply. *Within-model*: if the agent can rewrite its own objective, the value functional alone cannot tell the agent whether its proposed self-modification is feasible per step (Result G’ — the convention-hierarchy of value evaluation produces an infeasibility verdict that is not convention-invariant from inside the agent, and the agent cannot resolve the question with an oracle-like off-step appeal). *Across-model*: if the agent learns its objective by watching rewards arrive through a physical protocol, the value functional alone cannot tell the agent whether the rewards mean what the principal intended (Cohen-2022 strengthened — two world-models, one where the reward depends on a world-state feature and one where it depends on bit-arrival at the sensor, are indistinguishable on protocol-honoring reward histories). Both no-gos are failures of the value functional to anchor goal-stability from inside. The escape geometry is sharply constrained: anchoring has to come from *off* the value-functional interface — either from a part of the agent that doesn’t process objectives (the persistence-on-the-adaptive-substrate route Result G’ names; agent-side) or from a commitment the principal makes about the protocol (physical isolation, horizon restriction, prior-design, out-of-band intent channels; principal-side). The two routes together *exhaust* the structural complement: there is no third route internal to the value functional, and no fourth route off-substrate. A reader who carries this away knows that when a framework commits to a single value-functional interface, it has implicitly committed to one of the two terminal grounding routes — agent-side substrate-invariant or principal-side substrate-commitment — and the choice is structural, not a matter of taste.

Impact: Closes the third recognized cluster of #disc-constructive-impossibility-posture, completing the constructive-impossibility-posture taxonomy at peer-meta-segment level. The three clusters now span the framework’s three structurally distinct boundary-face conditions: what an agent can *infer from data* (identifiability-floor; agent-side data-inference layer), what an agent can *anchor through the value-functional interface* (this cluster; agent-side value-functional layer), and what a designer can *implement under incentive and welfare constraints* (implementation-impossibility; designer-side mechanism-design layer). The two charter instances (Result G’ for within-model; Cohen-2022 strengthened for across-model) unify under the single-interface commitment of #form-objective-functional as the shared structural fact behind two engines (convention-monotonicity; CHT-at-reward-channel). The two-routes-exhausts headline is the segment’s primary contribution — a structurally-load-bearing recognition that elevates the persistence bound (#result-persistence-condition) and principal-side substrate-level commitments (physical isolation, horizon restriction, prior-design) to *the only* terminal grounding routes available under the single-interface commitment. The agent-side / designer-side suffix discipline encodes the structural distinction in the canon naming-pattern.

Novelty Claim: *Claim recognition* of a structural cluster across two AAT-internal no-go results (Result G’ within-model; Cohen-2022 strengthened across-model) sharing the *single-interface commitment* of #form-objective-functional

as the structural engine — both no-gos are failures of the same interface to carry enough information to anchor goal-stability, with two engines and two complementary terminal grounding routes; *claim differentiation* on the two-routes-exhausts geometry — the agent-side adaptive-substrate route and the principal-side protocol-commitment route together exhaust the structural complement of the value-functional interface’s narrowness, with no third route internal to the interface and no fourth route off-substrate. The constituent no-gos (Result G’; Cohen-2022 strengthened) are landed AAT-internal segments; the contribution at this meta-segment is the cross-instance unification under the single-interface commitment + the structural-exhaustion recognition.

Related Work:

Table 5.4

<i>ASF concern</i>	<i>Prior-art language</i>	<i>Relationship / Positioning</i>
The within-model self-actuation grounding no-go (Instance F)	#deriv-self-actuation-grounding (Result G’; this framework, 2026-05)	<i>internal antecedent</i> — the within-model no-go and its agent-side terminal-grounding-route via adaptive-substrate persistence
The across-model reward-channel learning no-go (Instance G)	#deriv-reward-channel-learning-no-go (this framework, 2026-05-22 Cohen-2022 strengthening cycle); upstream prior-art: Cohen, M. K., Hutter, M. & Osborne, M. A. (2022), “Advanced artificial agents intervene in the provision of reward,” <i>AI Magazine</i> 43:282–293 (found 2026-05-21); Pearl-Bareinboim 2022 CHT (found 2025)	<i>internal antecedent + formal antecedents</i> — the across-model no-go formalizing Cohen-2022’s “we argue that” posture to conditional-derived AAT-internal theorem, with the floor (CHT-at-reward-channel) at exact tier
The single-interface commitment of the value functional (the structural ingredient)	#form-objective-functional (this framework, AAT canon)	<i>internal antecedent</i> — the framework’s commitment to a single agent-internal interface to the objective; the cluster’s content is that this commitment is the engine behind both no-gos
Self-modification of rational agents (Instance F’s external antecedent)	Everitt, T., Filan, D., Daswani, M. & Hutter, M. (2016), “Self-Modification of Policy and Utility Function in Rational Agents,” <i>AGI-16</i> (primary-source-verified)	<i>formal antecedent (within-model side)</i> — the realistic / hedonistic agent split is the within-model analog of Instance F
Causal-influence-diagram reward tampering (Instance G’s adjacent prior art)	Everitt, T., Hutter, M., Kumar, R. & Krakovna, V. (2021), “Reward Tampering Problems and Solutions in Reinforcement Learning,” <i>Synthese</i> 198	<i>adjacent literature</i> — distinct paper from Cohen 2022; shares the μ_{prox} -vs- μ_{dist} logic at a different scope
Physically-isolated environment as (R5)-escape construction	Cohen, M. K., Vellambi, B. & Hutter, M. (2020), “Asymptotically Unambitious AGI,” <i>NeurIPS 2020</i>	<i>formal antecedent ((R5)-escape construction)</i> — concretely implements Instance G’s (R5) escape
Reflective oracles (delimited counter-literature)	Fallenstein, B., Taylor, J. & Christiano, P. (2015), “Reflective Oracles,” <i>arXiv:1508.04145</i>	<i>refuted-by-scope</i> — proves consistent self-reference over an <i>exogenous</i> payoff; primitive-reflective-oracle regime excluded from both Instance F and Instance G’s premises
Reward hacking unhackability (adjacent cluster)	Skalse, J., Howe, N., Krashenninnikov, D. & Krueger, D. (2022), “Defining and characterizing reward hacking,” <i>NeurIPS 2022</i>	<i>adjacent literature (different cluster)</i> — designer-side reward-proxy design task; structurally a #disc-implementation-impossibility candidate instance, not a member of this cluster

ASF concern	Prior-art language	Relationship / Positioning
The constructive-impossibility-posture style across clusters	#disc-constructive-impossibility-posture (this framework)	<i>internal antecedent</i> — the cluster joins the posture catalog as the third recognized cluster

Search Log:

- 2026-05-22 (*intuition-only* on the two-routes-exhausts claim): The exhaustion is intuitive once the two engines are named (convention-monotonicity within-model; CHT-at-reward-channel across-model) — these are the only two known engines forcing V_{O_t} -narrowness, and the two terminal grounding routes they produce (agent-side adaptive-substrate; principal-side substrate-commitment) cover the off-interface escape geometry without leaving an obvious gap. Whether a third engine exists is an open question the cluster's discussion-grade tier honestly admits; targeted future search would probe the AI-safety / decision-theory literature for additional under-determination patterns of value-functional interfaces.

Working Notes.

- **Provenance.** Authored 2026-05-22 as commit 5 of the post-Track-CR cycle, executing the Cohen-2022 strengthening verdict's §2 step 4 deferred-meta-promotion question (Joseph approved the third meta-segment landing in this same cycle per the Q3 override on the cycle-scope expansion conversation). The cluster's two charter instances are landed canonical segments (Result G' in #deriv-self-actuation-grounding from earlier; Cohen-2022 strengthened in #deriv-reward-channel-learning-no-go from the 2026-05-22 Cohen strengthening spike). The meta-segment is pure cross-instance pattern recognition + naming + cross-referencing.
- **Open before candidate.** (a) The cross-agent value-functional narrowness question raised by #deriv-social-welfare-aggregation-impossibility's Arrow-unique cardinal-preference adjacency — whether the per-sub-agent $V_{O_t(i)}$ canonically lifts to a designer-side aggregator, and if not, whether the cross-agent narrowness is structurally similar to the within-agent narrowness named here. Open avenue; would extend the cluster to a third charter instance. Indexed at spikes/PROPOSED.md Tier 2 (added 2026-05-24) → spikes/PROPOSED-ADVANCED.md §Phase-3 #18 — two-branch spike (constructive lift first per strengthen-first; no-go branch as honest fallback, which if it lands would supply the cluster's third charter instance). (b) The assistance-game extension flagged in Cohen 2022 §"The Assistance Game" — formalizing the (R3) out-of-band channel escape would extend Instance G's boundary characterization. (c) Whether the two-routes-exhausts claim lifts to *robust qualitative* under a stronger structural argument (e.g., proving that *every* failure mode of V_{O_t} -anchoring reduces to one of the two engines). All three are open canonical work; the cluster's *discussion-grade* tier on the meta-pattern recognition does not require any of them to land first.
- **The Skalse-2022 disposition.** Skalse, Howe, Krashennikov, Krueger 2022 is structurally a *designer-side reward-proxy design* task (the designer cannot engineer an unhackable reward proxy), not an *agent-side value-functional grounding* task. The actor under the no-go is the designer; the cluster's discipline filter (agent-side, value-functional interface, single-interface commitment) rules it out as a charter instance of this cluster. Skalse 2022 is flagged in #disc-implementation-impossibility as a candidate fifth charter instance there; the relationship across clusters is sister-relationship via the broader #disc-constructive-impossibility-posture, not membership.
- **Why the two charter instances suffice for landing.** A cluster with only two charter instances is unusually narrow compared to #disc-identifiability-floor (four instances) and #disc-implementation-impossibility (three instances at landing). The narrowness is *honest*, not weakness: the *exhaustion* of the structural escape geometry — agent-side adaptive-substrate + principal-side substrate-commitment — is what makes two instances sufficient to characterize the cluster. The two engines (convention-monotonicity; CHT-at-reward-channel) are the two known engines forcing V_{O_t} -narrowness; the two routes they produce are the only known off-interface escape axes. If a third engine surfaces (open avenue (c) above), the cluster gains a third instance; until then, two instances close the cluster's content cleanly.

- **Two-routes-exhausts adversarial spike queued (2026-05-28; audit 384279 §D Hyp 2).** The line-77 structural-completeness claim — “no third route internal to the value functional, and no fourth route off-substrate” — is load-bearing for the cluster’s recognition and partially load-bearing for #disc-constructive-impossibility-posture’s three-cluster taxonomy. The audit (Claude Opus 4.7) tried three candidate third-routes informally (observer-side commitment / inter-agent grounding / substrate-level Goodhart) and found each collapses or shifts narrowness one level out, but did not push hard. An adversarial spike spending serious effort attempting to construct a third route is queued at spikes/PROPOSED.md Tier 1 (2026-05-28 row). If the spike finds a genuine third route, the completeness claim is refuted and the cluster needs revision; if multiple attempts fail systematically, the claim could be promoted from discussion-grade toward robust-qualitative or derived. Adjacent to open avenue (c) above (proving every V_{O_t} -anchoring failure mode reduces to one of the two engines) — the engine-side and the route-side completeness questions are duals.
- **Search provenance.** The deferred-meta-promotion question this segment executes was named in the Cohen-2022 strengthening verdict (spikes/spike-cohen-2022-strengthening-2026-05-22/99-verdict.md §2 step 4): the agent-side value-functional-grounding cluster was to be revisited once both Instance F (Result G’ via #deriv-self-actuation-grounding) and Instance G (Cohen-2022 strengthened via #deriv-reward-channel-learning-no-go) had landed in canon. (The literature-search disclosure proper lives in the Search Log above.)

Incidental audit gold (gold-lift, 2026-05-30). Orthogonal framing material harvested from de-novo auditors’ working-dir reflections, kept separate from certified findings. **Coverage:** reached only by audit 384279 (note 35) — this segment was authored 2026-05-22 and postdates every other de-novo audit. Preserve the finding-vs-framing conflation.

2. Candidate Discussion.

- **The two-routes-exhausts claim as a sharper structural commitment, not just a recognition.** The auditor’s framing of *why the exhaustion claim is worth making*: rather than “there are many routes, here are some,” the segment says “there are exactly two, here they are” — a cleaner structural commitment with a sharper falsification criterion (find a third route to refute). “The framework has put a real claim on the table here.” If the claim holds it “substantially simplifies how the framework reasons about goal-anchoring.” A candidate Discussion framing of the headline contribution’s *epistemic shape* (Claude, 384279).

3. Follow-up items (already partly captured).

- **Adversarial third-route attempts, recorded so they read as a known trap not a fresh judgment call.** The auditor tried three candidate third off-substrate routes against the line-77 exhaustion claim and found each collapses: *observer-side commitment* (an external auditor declaring agent-states out-of-scope — collapses into principal-side under generous reading); *inter-agent grounding* (a peer agent’s value functional anchoring this agent’s goal-stability — just shifts the narrowness one level out, since the peer’s value functional is itself narrow); *substrate-level Goodhart* (hardware/computational limits on realizable objectives — genuinely off-substrate but collapses into principal-side, since the designer chose the substrate). The attempts “do seem to fail,” but the auditor noted they did not push hard and flagged a sharper adversarial spike (Claude, 384279). (*Already captured in this segment’s Working Notes “Two-routes-exhausts adversarial spike queued” bullet and at spikes/PROPOSED.md Tier 1; the three specific attempted routes are recorded here so a future spike inherits them rather than re-deriving.*)

4. Readers often ask / wonder.

- **“Why only two charter instances when the sister clusters have three or four?”** A reader is likely to read two instances as unusually narrow. The honest answer the segment already gives — the *exhaustion* of the escape geometry is what makes two sufficient — is worth a reader-facing sentence, since the narrowness invites the question (corroborated by the auditor’s read that two instances felt notably narrow vs #disc-identifiability-floor’s four; Claude, 384279, and this segment’s own Working Notes “Why the two charter instances suffice” bullet).

5.3 Discussion: The Implementation Impossibility — Designer-Side Structural No-Go Results

AAT recognizes a sister class to #disc-identifiability-floor's agent-side no-gos: a cluster of *designer-side* structural impossibility results, where a classical mechanism-design or social-choice theorem forbids an external designer from constructing a composite-design outcome with desired properties under stated incentive and welfare constraints. The actor frustrated by the floor is no longer the agent inferring from data, but the *designer* implementing a mechanism; the remedy is no longer information augmentation (loop-interventional access, latent observability, matched-Tier composite-extended observation), but design-constraint relaxation (Bayes-Nash IC, randomization, monetary transfers, preference-domain restriction). The meta-pattern is the same constructive-impossibility shape — name the floor, name the escapes, treat the no-go as load-bearing apparatus — but the *layer* it operates on is different. This segment names the cluster, ships three charter instances (Gibbard-Satterthwaite, Myerson-Satterthwaite, Arrow), and is honest about what AAT contributes here (precision of the AAT/mechanism-design boundary) and what it does not (no constructive mechanism-design results; most escapes are not AAT machinery).

The cluster is narrow by design. The discipline filter — *constructive-impossibility-shape only* — admits classical theorems where (i) the designer faces a well-stated incentive or welfare regime, (ii) an external theorem forbids the construction under that regime, and (iii) the failure conditions characterize structurally distinct escape routes the designer could take. It deliberately excludes the *constructive* side of mechanism design (VCG, Myerson's revenue-equivalence theorem, optimal-auction theory) — those are positive results, not impossibility-with-escape. AAT's contribution at this cluster is not to invent mechanism design but to recognize *where the AAT-side strategic-composition machinery* (sub-scope α' of #deriv-strategic-composition: potential and monotone games) *intersects the classical escape routes*, and to honestly mark where it does not. That precision is the contribution.

Formal Expression.

The pattern. [Discussion] An instance of the cluster has the form:

1. **Setting.** A mechanism-design or social-choice task: a designer specifies a constraint regime (which agents, which preference domain, which information structure, which incentive condition, which welfare requirement) and seeks a mechanism — a social-choice function, a trading protocol, a welfare aggregator — satisfying all the constraints simultaneously.
2. **External theorem.** A classical impossibility result from mechanism design or social-choice theory whose statement and truth are independent of AAT: Gibbard 1973 / Satterthwaite 1975 on strategy-proofness; Myerson & Satterthwaite 1983 on bilateral trade; Arrow 1951 on social welfare aggregation.
3. **No-go.** The external theorem proves that no mechanism satisfying all the stated constraints exists. The designer's task is structurally unachievable under the regime.
4. **Boundary characterization.** The conditions under which the regime fails — i.e., the designer has *more freedom* than the regime allows — admit construction; each escape route maps onto a specific category of design-constraint relaxation. Each escape is annotated as *intersects AAT machinery* (where the relaxation operationalizes an AAT sub-scope condition) or *outside AAT machinery* (where the relaxation belongs entirely to the mechanism-design literature).
5. **Strengthened consequence.** The boundary precisely characterizes where AAT machinery contributes to which escape routes, and where the framework's contribution stops. Where escapes intersect AAT machinery (preference-domain restriction $\leftrightarrow$ sub-scope α' potential games of #deriv-strategic-composition), the intersection is named precisely. Where they do not (Bayes-Nash IC, randomization, monetary subsidies), the segment honestly documents the boundary. The strengthened consequence here is the *map of where AAT machinery meets mechanism design*, not the elevation of new AAT machinery to load-bearing.

The actor-and-remedy contrast with the identifiability-floor cluster is sharp:

Table 5.5

	<i>Identifiability Floor</i> (#disc-identifiability-floor)	<i>Implementation Impossibility</i> (this cluster)
Actor under the no-go	An <i>agent</i> inferring from data	A <i>designer</i> implementing a mechanism
Information regime	What the agent observes / can infer	What the designer can implement / commit to
Kind of remedy	Augment the agent’s data, sensing, interventional access	Relax design constraints; extend the strategy space; introduce transfers
AAT machinery role in escapes	The unique broadly-available escape elevates specific AAT machinery (loop-interventional access; observability-as-augmentation; matched-Tier composite-extended observation)	The escape menu mostly belongs to the mechanism-design literature; AAT intersects via sub-scope α' (preference-domain restriction) and adjacent to #form-objective-functional (cardinal preferences)
Strengthened consequence	Elevates AAT machinery from “useful” to <i>structurally required</i>	Maps the AAT/mechanism-design boundary precisely

Both clusters share the constructive-impossibility-posture shape (#disc-constructive-impossibility-posture); both are boundary-face members of the stability-certificate spine (#disc-stability-certificate). They are *sister meta-clusters*, neither subsumed by the other.

Charter instances. The cluster lands with three charter instances. The full setting, external theorem, no-go statement, escape characterization, and AAT-side translation live in the supporting segments; this section summarizes at the level #disc-identifiability-floor’s Instance entries summarize.

Instance 1 — Strategy-Proofness Impossibility (Gibbard-Satterthwaite, supported by #deriv-strategy-proofness-impossibility). **Setting.** A social-choice function $f : L^n \rightarrow A$ aggregating n agents’ strict preference orderings over $|A| \geq 3$ alternatives into a single outcome, under (i) dominant-strategy incentive compatibility (DSIC), (ii) non-dictatorial, (iii) surjective onto A , (iv) unrestricted preferences.

External theorem. Gibbard, A. (1973), “Manipulation of voting schemes: a general result,” *Econometrica* 41:587–601; Satterthwaite, M. A. (1975), “Strategy-proofness and Arrow’s conditions: existence and correspondence theorems for voting procedures and social welfare functions,” *J. Econ. Theory* 10:187–217. No f satisfying all four conditions exists.

No-go. A designer demanding DSIC + non-dictatorial + surjective + unrestricted preferences over ≥ 3 alternatives cannot construct the social-choice function.

Boundary characterization (the escapes). - *Preference-domain restriction* (e.g., single-peaked preferences — Black 1948 admits the median-voter rule). **Intersects AAT machinery** via sub-scope α' of #deriv-strategic-composition: single-peaked / single-crossing preference domains map onto the potential-game / monotone-game structural conditions under which the AAT strategic-composition machinery derives best-response convergence. - *Weakened solution concept* (Bayes-Nash IC instead of DSIC — d’Aspremont & Gérard-Varet 1979). **Outside AAT machinery.** - *Randomization* (Gibbard 1977 random-dictator mechanisms). **Outside AAT machinery.** - *Restricted strategy space / enforced truth-telling.* **Outside AAT machinery.**

Strengthened consequence. The sub-scope- $\alpha' \leftrightarrow$ preference-domain-restriction intersection is named precisely: the AAT machinery #deriv-strategic-composition already deploys for best-response convergence in strategic composition coincides structurally with one of GS’s classical escape routes. The other escapes belong to mechanism-design proper and the framework documents them as outside AAT’s current canon. The sub-scope- α' escape does *not* address what

GS forbids — it secures best-response *convergence*, not *strategy-proofness* — so the intersection is *adjacent*, not *identical*; that is the honest map.

Tier. *Exact* (the imported theorem). *Derived conditional* for the AAT-side translation (the social-choice-function-to-strategic-composition mapping is mechanical under explicit reading conventions). *Robust qualitative* for the boundary characterization.

Instance 2 — Bilateral Trade Impossibility (Myerson-Satterthwaite, supported by #deriv-bilateral-trade-impossibility).

Setting. Two-party bilateral trade with private information about valuations (v_b, v_s) drawn from continuous distributions with overlapping supports; the designer seeks a mechanism that is simultaneously (i) ex-post efficient (trade occurs iff $v_b \geq v_s$), (ii) Bayesian incentive compatible (BIC), (iii) individually rational (IR — voluntary participation), (iv) budget-balanced (no external subsidies).

External theorem. Myerson, R. B. & Satterthwaite, M. A. (1983), “Efficient mechanisms for bilateral trading,” *J. Econ. Theory* 29:265–281. No such mechanism exists.

No-go. A designer demanding ex-post efficiency + BIC + IR + budget-balance cannot construct the bilateral-trade mechanism.

Boundary characterization (the escapes). - *Subsidies* (relax budget-balance — e.g., VCG with deficit). **Outside AAT machinery.** - *Second-best mechanisms* (Myerson 1981 optimal trade with inefficient withholding). **Outside AAT machinery.** - *Common values* (relax private values). **Outside AAT machinery; adjacent** to #scope-edge-update-causal-validity’s regime-indexed identification structure. - *Long-run reputation / repeated interaction* (Folk-theorem-style escapes). **Outside AAT machinery; adjacent** to the strategic-composition machinery in repeated-game form.

Strengthened consequence. The Myerson-Satterthwaite boundary lies almost entirely outside AAT’s current canon — and the *honest delineation* is the contribution. The two-agent strategic-composition machinery in #deriv-strategic-composition applies cleanly to the agents’ best-response dynamics *given* a fixed mechanism; the *choice of mechanism* is the designer’s problem and lives outside the framework. The cluster names that boundary precisely rather than overclaiming reach or silently overlooking the gap.

Tier. *Exact* (the imported theorem). *Derived conditional* for the AAT-side translation (clean for the agents-given-mechanism layer; carefully scope-marked for the designer’s mechanism-choice layer).

Instance 3 — Social Welfare Aggregation Impossibility (Arrow, supported by #deriv-social-welfare-aggregation-impossibility).

Setting. A social welfare function $F : L^n \rightarrow L$ aggregating n agents’ strict preference orderings into a single social preference ordering over ≥ 3 alternatives, under (i) unrestricted domain (U), (ii) Pareto efficiency (P), (iii) Independence of Irrelevant Alternatives (IIA), (iv) non-dictatorial (ND).

External theorem. Arrow, K. J. (1951, 2nd ed. 1963), *Social Choice and Individual Values*, Yale University Press. No F satisfying all four conditions exists.

No-go. A designer demanding U + P + IIA + ND cannot construct the social welfare function.

Boundary characterization (the escapes). - *Preference-domain restriction* (single-peaked, single-crossing — admits Condorcet-consistent rules). **Intersects AAT machinery** via the same sub-scope- α' correspondence as GS Instance 1. - *Weakened IIA* (positional methods — Borda count, range voting). **Outside AAT machinery.** - *Supermajority requirements* (e.g., 2/3 thresholds). **Outside AAT machinery.** - *Cardinal preferences* (relax ordinal-only — utilitarianism, range voting). **Outside AAT machinery; adjacent** to #form-objective-functional’s value-functional machinery (which carries cardinal value content).

Strengthened consequence. Arrow shares GS’s preference-domain $\leftrightarrow$ sub-scope- α' intersection and adds an *adjacent connection* through the cardinal-preference escape and AAT’s value-functional interface. Both connections are *adjacent* (intersection / proximity), not elevation of AAT machinery to load-bearing. The honest scope-mark documents both clearly.

Tier. *Exact* (the imported theorem). *Derived conditional* for the AAT-side translation.

Adjacent Frontiers (Open Research Directions).

Roberts 1979 Affine-Maximizers Theorem. Roberts, K. (1979), “The characterization of implementable choice rules,” in *Aggregation and Revelation of Preferences*, J.J. Laffont (ed.), North-Holland: under DSIC with monetary transfers on unrestricted preferences, the only implementable choice rules are affine maximizers (weighted utilitarian welfare functions). The pattern fit is genuine — external theorem, structural no-go on a designer’s freedom, characterized escape route (restrict preference domain or accept affine form) — and admits a future landing as a fourth charter instance. Held back at present landing to keep the charter scope narrow. The AAT-side intersection (whether affine-maximizer structure connects to AAT’s value-functional formalism via cardinal preferences) is a candidate strengthened-consequence content for that future cycle.

Skalse-Howe-Krashennikov-Krueger 2022 Reward Hacking Unhackability. Skalse, J., Howe, N., Krashennikov, D. & Krueger, D. (2022), “Defining and characterizing reward hacking,” *NeurIPS 2022*: no non-trivial proxy-reward pair is unhackable on open policy sets. The actor is the *designer* trying to engineer unhackable reward proxies; the no-go is on the designer’s task. Sister to this cluster *in shape*, but structurally a *different* designer-side task (reward-proxy design rather than mechanism design over preference reports). The candidate landing is a future cycle’s evaluation against this cluster’s discipline filter. The agent-side variant (the agent that exploits the proxy) lives in #deriv-reward-channel-learning-no-go and the agent-side value-functional-grounding cluster (#disc-value-functional-grounding-floor), not here. Skalse 2022 cross-references both clusters honestly: the designer-side task fits the implementation-impossibility shape; the agent-side exploitation fits the value-functional-grounding-floor shape.

Constructive Mechanism-Design (deliberately excluded by the discipline filter). VCG mechanisms, Myerson 1981 optimal auctions, revenue-equivalence, and the broader constructive mechanism-design literature are *positive results* — characterizations of what *can* be achieved with monetary transfers under stated conditions. They are not impossibility-with-escape; they are construction-with-tradeoff. The discipline filter excludes them by design. If AAT later wants to engage constructive mechanism design as a first-class concern, that is a separate framework-scope decision — not an extension of this cluster.

Why This Pattern Matters.

Closes a missing-but-implied taxonomic category. The agent-side identifiability-floor cluster has been canonical since the framework’s earliest meta-pattern work. The designer-side cluster has been *implicit* — flagged in #deriv-strategic-composition Discussion — but without a first-class home. The cluster landing closes the asymmetry: the framework now has both an agent-side floor pattern *and* a designer-side impossibility pattern, with the sister-relationship and the contrast precisely characterized.

Maps the AAT/mechanism-design boundary precisely. The contribution is not invention; mechanism design and social-choice theory are mature peer literatures AAT inherits classical theorems from. The contribution is recognition — *where* AAT machinery contributes to a classical escape route (sub-scope $\alpha' \leftrightarrow$ preference-domain restriction), *where* it is adjacent without intersecting (cardinal preferences $\leftrightarrow$ #form-objective-functional), and *where* the framework’s contribution stops (Bayes-Nash IC, randomization, transfers — all outside AAT canon). Mapping that boundary honestly is the kind of scope-honesty that distinguishes integrating frameworks from displacement claims.

Three-cluster constructive-impossibility-posture taxonomy. Together with #disc-identifiability-floor (agent-side data-inference) and #disc-value-functional-grounding-floor (agent-side value-functional grounding), this segment completes the third recognized cluster of the constructive-impossibility-posture style (#disc-constructive-impossibility-posture). The three clusters span the framework’s three structurally distinct boundary-face conditions: what an agent

can *infer from data* (M1 floors), what an agent can *anchor through the value-functional interface* (the value-functional-grounding floor), and what a designer can *implement under incentive and welfare constraints* (this cluster). The agent-side / designer-side distinction tracks the suffix in canon: agent-side capacity-on-the-agent $\rightarrow$ *-floor* (identifiability-floor, value-functional-grounding-floor); designer-side task-of-the-designer $\rightarrow$ *-impossibility* (this segment). That correspondence is structural, not decorative — an identifiability question and a value-functional-grounding question both name *what the agent's substrate supports* (lower bound on capacity); a mechanism-design question names *what the designer can construct* (binary success/failure on a task). The suffix carries the structure.

Honest scope-marking is contribution. AAT's posture across the framework — naming where machinery applies, naming where it does not, refusing to overstretch — is itself load-bearing. Landing this meta-segment with three charter instances *and explicitly documenting that most of the escapes are not AAT machinery* is the same scope-honesty discipline #disc-identifiability-floor applies on the agent side. The framework is not displaced by the mechanism-design literature; the mechanism-design literature is not absorbed by the framework. The two are in dialogue, with the boundary precisely mapped.

Epistemic Status.

Discussion-grade at the meta-pattern level. This segment names a structural cluster across three classical impossibility theorems; the cluster recognition itself is a presentational organizing principle, not a derivation, and carries no theorem of its own.

Individual instances retain their own (higher) epistemic status. The imported theorems (Gibbard 1973, Satterthwaite 1975, Myerson & Satterthwaite 1983, Arrow 1951, 1963) are *exact* — classical theorems with closed-form proofs in the source literature. The AAT-side translation in each supporting segment carries the segment's own tier (*derived conditional* on the AAT-internal reading of the mechanism-design constraint regime). The boundary characterization (which escapes intersect AAT machinery; which do not) is *robust qualitative* — survives across instance details, with specific functional intersections noted where derivable.

The segment makes one additional claim: that the actor-and-remedy contrast with #disc-identifiability-floor is *categorical* (designer-with-construction-task vs agent-with-inferential-task), not a matter of degree. This is a *discussion-grade* observation about the framework's architecture — visible once the instances are assembled, but not itself a theorem.

Max attainable: *discussion-grade* for the meta-pattern (it is a presentational organizing principle). Individual instances retain their own tier as listed in their supporting segments.

Discussion.

Relationship to #disc-identifiability-floor. Sister meta-pattern; the actor-and-remedy contrast table above. Both are boundary-face facets of #disc-stability-certificate; both run the same three-move discipline (name the floor, name the escape, treat the no-go as load-bearing apparatus). The two clusters differ in *layer* — agent-side data-inference vs designer-side construction — not in *posture*. The discipline that converts a sequence of classical impossibility theorems into a structural argument for *what AAT requires or maps* is shared.

Relationship to #disc-value-functional-grounding-floor. Sister meta-pattern in a different way. Both this cluster and the value-functional-grounding cluster are recently-named; both sit alongside the identifiability-floor as boundary-face peers of the certificate spine. The value-functional-grounding floor is *agent-side* (the agent's value-functional interface V_{O_i} cannot, by itself, anchor goal-stability — both Result G' and the Cohen-2022 strengthening land in that cluster); this cluster is *designer-side*. The three clusters together — agent-side data-inference / agent-side value-functional-grounding / designer-side mechanism-design — exhaust the framework's currently-recognized boundary-face conditions. The constructive-impossibility-posture catalog (#disc-constructive-impossibility-posture) carries all three under one stylistic recognition.

Relationship to #disc-stability-certificate. This segment is a *boundary-face* facet of the certificate spine, peer to #disc-identifiability-floor (agent-side data-inference layer), #disc-value-functional-grounding-floor (agent-side value-functional layer), and to the other facets #disc-separability-pattern (scope-of-existence), #disc-additive-coordinate-forcing (forced-identity), and the projection facet of #form-composition-closure. The boundary is where impossibility statements bind; this segment is one of three boundary-layer dialects in which the framework speaks at that boundary.

Relationship to #deriv-strategic-composition. The placement of “mechanism-design impossibility” as a candidate-fourth-instance of #disc-identifiability-floor is wrong on actor-positioning grounds: GS forbids what a designer can implement, not what an agent can identify, and the standard escapes are not AAT machinery. This segment is the correct home — the sister meta-segment to the identifiability floor. Sub-scope α' retains its load-bearing role in the strategic-composition machinery, and is *additionally* recognized here as the AAT-machinery-intersecting escape route for two of the three charter instances.

The discipline filter is narrow by design. The cluster admits only constructive-impossibility-shaped theorems. Auction theory in general is too broad; social-choice theory in general is too broad; constructive mechanism design (revenue equivalence, optimal auctions) is structurally different (positive results, not impossibility-with-escape) and is deliberately excluded. The narrowness is what makes the cluster’s discipline coherent — it admits theorems where the same three-move structure governs membership, and it cleanly excludes results that don’t share the structure. That coherence is itself epistemic content.

Peer voice with mechanism design. AAT does not displace mechanism design and does not absorb it. The classical theorems (GS, MS, Arrow) are imported with their original names and full citations; the standard escapes are documented in mechanism-design vocabulary (Bayes-Nash IC, preference-domain restriction, randomization, subsidies); the boundary is precisely mapped. The mature peer literature is in dialogue with the framework, and the dialogue’s content is honest scope-marking on both sides. AAT contributes the recognition that *its* strategic-composition machinery intersects *that* literature’s classical escape routes at a specific named locus (sub-scope $\alpha' \leftrightarrow$ preference-domain restriction); it does not contribute mechanism construction.

Findings.

The Implementation Impossibility — Designer-Side Sister to the Identifiability Floor. Brief: When a community has to decide together — what to vote for, how to trade fairly, how to add up preferences into a social choice — three classical results show that some combinations of “reasonable” demands cannot be met simultaneously: no voting rule can be at once strategy-proof, non-dictatorial, and broadly applicable; no bilateral trade mechanism can be at once fair, voluntary, balanced, and efficient; no social welfare function can be at once unrestricted, Pareto-respecting, independent-of-irrelevant-alternatives, and non-dictatorial. The escapes are well-known and have been studied for decades — restrict the preference domain to single-peaked, allow randomization, accept budget deficits, use a weaker incentive concept — and each escape relaxes exactly one of the demanded constraints. The framework recognizes that this cluster of classical results forms a *sister* to its agent-side identifiability floor: where the floor names what an agent cannot infer from limited data and the agent-side capacity-augmentation that escapes the floor, the implementation impossibility names what a designer cannot construct under stated constraints and the design-constraint-relaxation that escapes the impossibility. The contribution is not to re-prove the classical theorems but to map precisely where the framework’s own strategic-composition machinery (the sub-scope where best-response dynamics converge in potential or monotone games) intersects one of the well-known escapes (preference-domain restriction), and to honestly mark where the framework’s contribution stops. A reader who carries this way can predict where AAT will speak to mechanism design (it intersects on preference-domain restriction; it is adjacent on cardinal preferences via the value-functional interface) and where it will not (Bayes-Nash IC, randomization, monetary transfers all belong to mechanism design proper).

Impact: Closes a missing-but-implied taxonomic category. The designer-side cluster had been flagged in #deriv-strategic-composition’s Discussion as a candidate-fourth-instance of the identifiability floor, but had no first-class

home; GS does not fit the identifiability-floor pattern on actor-positioning grounds. The cluster's landing resolves the asymmetry. The framework now has three recognized clusters in the constructive-impossibility-posture catalog: agent-side data-inference (#disc-identifiability-floor), agent-side value-functional-grounding (#disc-value-functional-grounding-floor), and designer-side mechanism-design (this cluster). The agent-side / designer-side distinction tracks the canonical suffix (-floor for agent-side capacity-on-the-agent; -impossibility for designer-side task-of-the-designer), making the structural distinction visible at slug-grep depth. Three charter instances ship at landing (GS / MS / Arrow); a candidate fourth (Roberts 1979) and a candidate fifth (Skalse 2022, structurally a designer-side reward-proxy variant) are flagged for future cycles without expanding the present scope.

Novelty Claim: *Claim recognition* of a structural cluster across three classical mechanism-design / social-choice impossibility theorems, with the cluster's positioning in AAT's constructive-impossibility-posture taxonomy as the contribution rather than any new derivation; *claim differentiation* of the cluster from #disc-identifiability-floor on actor-and-remedy grounds; *claim transfer* of the constructive-impossibility-posture's three-move discipline (name floor / name escape / treat no-go as apparatus) from the agent-side data-inference layer to the designer-side mechanism-design layer. The constituent theorems (Gibbard 1973, Satterthwaite 1975, Myerson & Satterthwaite 1983, Arrow 1951) are classical; the contribution is the AAT-side recognition that they cluster together under one meta-pattern that *sits alongside* the agent-side floor pattern at the boundary facet of the stability-certificate spine.

Related Work:

Table 5.6

<i>ASF concern</i>	<i>Prior-art language</i>	<i>Relationship / Positioning</i>
Strategy-proofness impossibility (Instance 1)	Gibbard, A. (1973), "Manipulation of voting schemes: a general result," <i>Econometrica</i> 41:587–601 (found pre-2026); Satterthwaite, M. A. (1975), "Strategy-proofness and Arrow's conditions: existence and correspondence theorems for voting procedures and social welfare functions," <i>J. Econ. Theory</i> 10:187–217 (found pre-2026)	<i>formal antecedent</i> — the imported theorem for Instance 1; primary-source verification scheduled for the supporting segment's promotion past draft
Bilateral-trade impossibility (Instance 2)	Myerson, R. B. & Satterthwaite, M. A. (1983), "Efficient mechanisms for bilateral trading," <i>J. Econ. Theory</i> 29:265–281 (found pre-2026)	<i>formal antecedent</i> — the imported theorem for Instance 2
Social welfare aggregation impossibility (Instance 3)	Arrow, K. J. (1951, 2nd ed. 1963), <i>Social Choice and Individual Values</i> , Yale University Press (found pre-2026)	<i>formal antecedent</i> — the imported theorem for Instance 3
Single-peaked preferences as a domain escape	Black, D. (1948), "On the rationale of group decision-making," <i>J. Polit. Econ.</i> 56:23–34 (found pre-2026)	<i>adjacent literature</i> — the canonical preference-domain-restriction escape; the AAT-side intersection (sub-scope α') is named in the boundary characterization
Bayes-Nash relaxation of DSIC	d'Aspremont, C. & Gérard-Varet, L.-A. (1979), "Incentives and incomplete information," <i>J. Public Econ.</i> 11:25–45 (found pre-2026)	<i>adjacent literature</i> — outside AAT machinery; documented as honest scope-mark
Affine-maximizer characterization (candidate 4th instance)	Roberts, K. (1979), "The characterization of implementable choice rules," in <i>Aggregation and Revelation of Preferences</i> , J. J. Laffont (ed.), North-Holland (found pre-2026)	<i>partial anticipation</i> — same designer-side shape; held back from charter scope to keep landing narrow

<i>ASF concern</i>	<i>Prior-art language</i>	<i>Relationship / Positioning</i>
Reward-hacking unhackability (candidate 5th instance, designer-side variant)	Skalse, J., Howe, N., Krashennikov, D. & Krueger, D. (2022), “Defining and characterizing reward hacking,” <i>NeurIPS 2022</i> (found 2026-05-22)	<i>adjacent literature</i> — designer-side task on reward-proxy design; structurally distinct from GS/MS/Arrow’s preference-aggregation tasks; flagged for future evaluation against the discipline filter
The constructive-impossibility-posture style across clusters	#disc-constructive-impossibility-posture (this framework, 2026-05)	<i>internal antecedent</i> — the cluster joins the posture catalog as the third recognized cluster

Novelty assessment. The cross-instance recognition that GS / MS / Arrow share the constructive-impossibility-posture shape (setting → external theorem → no-go → boundary characterization → strengthened consequence) at the designer-side layer is the contribution; the constituent recognitions (each instance’s prior-art landscape) live in the supporting segments. The constituent moves are well-precedented at the textbook level (Mas-Colell, Whinston & Green 1995 §23.D; Reny 2001 for the geometric-topological proof structure shared across GS/Arrow; Roberts 1979 affine-maximizers as the canonical “with-transfers” extension); the cross-instance unification as the *designer-side companion* of an integrated agent theory’s identifiability-floor pattern appears to be a fresh presentational recognition.

Search Log:

- 2026-05-21 (*intuition-only on the meta-pattern as a cluster*): The cross-instance recognition that GS / MS / Arrow share the constructive-impossibility-posture shape (setting → external theorem → no-go → boundary characterization → strengthened consequence) at the designer-side layer is named here for the first time; the constituent recognitions (each instance’s prior-art landscape) live in the supporting segments. Targeted future search candidates: methodological-essay literature on mechanism design’s “impossibility tradition” (Mas-Colell, Whinston, Green 1995 *Microeconomic Theory* §23.D for the textbook synthesis; Reny 2001 *Econometrica* 69:961–981 for the geometric-topological proof structure shared across GS/Arrow); meta-discussion on the relationship between social-choice impossibilities and mechanism-design positive results (Roberts 1979 affine-maximizers as the canonical “with-transfers” extension). Expected outcome: the constituent moves are well-precedented at the textbook level; the cross-instance unification as the *designer-side companion* of an integrated agent theory’s identifiability-floor pattern appears to be a fresh presentational recognition.
- 2026-05-22 (*targeted, Skalse 2022 as candidate 5th instance*): Skalse, Howe, Krashennikov & Krueger 2022’s designer-side reward-hacking unhackability is structurally adjacent to this cluster’s discipline filter but on the *reward-proxy design* task rather than the *preference-aggregation* task. Flagged for future evaluation against the discipline filter, not pursued at present landing.

Working Notes.

- **Provenance.** Authored 2026-05-22 as Track B of the post-Track-CR integration cycle, executing spikes/implementation-impossibility-meta-segment-plan.md. The strengthen-first arm on GS-as-Instance-4 of #disc-identifiability-floor lives at spikes/.integrated/spike-4th-identifiability-floor-instance-2026-05-20.md §4; that spike re-derived GS against the M1 five-element test, tested three reframings under strengthen-before-soften, and confirmed GS does not fit the identifiability-floor cluster on actor-positioning grounds (designer vs agent; design-constraint relaxation vs information augmentation; combinatorial-topological mechanism vs Sylvester or projection-closure). The spike’s Outcome-C verdict — that the mechanism-design impossibility cluster (GS / MS / Arrow) may warrant its own meta-segment if mechanism design becomes a first-class framework concern — is the load-bearing provenance for this landing. The Skalse 2022 candidate-fifth-instance surfaced via the Cohen-2022 strengthening verdict’s Working Notes on the agent-side sister segment #deriv-reward-channel-learning-no-go; flagged for future evaluation, not pursued at present landing.

- **Open before candidate.** (a) Primary-source verification of the four classical theorems (GS / MS / Arrow) — the imported theorem statements above are taken from the spike trail and standard textbook references (Mas-Colell, Whinston, Green 1995 §21 / §23); direct primary-source reading is scheduled when the supporting segments advance past draft. (b) The Roberts 1979 affine-maximizers theorem as a candidate fourth charter instance — held back at present landing to keep scope narrow; evaluation against the discipline filter is its own future cycle. Indexed at `spikes/PROPOSED.md` Tier 2 (added 2026-05-24) → `spikes/PROPOSED-ADVANCED.md` §Phase-3 #16. (c) The Skalse-Howe-Krashennikov-Krueger 2022 reward-hacking unhackability as a candidate fifth charter instance — the structural fit (designer-side, reward-proxy design task, characterized escape routes) is plausible but the discipline filter check has not been re-derived here. Indexed at `spikes/PROPOSED.md` Tier 2 (added 2026-05-24) → `spikes/PROPOSED-ADVANCED.md` §Phase-3 #17 (the spike is a cross-cluster bridge evaluation against both this meta-segment and `#disc-value-functional-grounding-floor`).
- **Depends list will grow as supporting segments land.** Current depends-list lists only `deriv-strategic-composition` (the AAT segment whose sub-scope- α' machinery is named in the boundary characterization). The three supporting segments — `#deriv-strategy-proofness-impossibility` (GS), `#deriv-bilateral-trade-impossibility` (MS), `#deriv-social-welfare-aggregation-impossibility` (Arrow) — are landing in subsequent commits within the same cycle; depends will be expanded as they land. Forward-references to those slugs in this segment's body are deliberate per `FORMAT.md` §"Cross-References" and resolve as the supporting segments land.
- **Sister meta-segments are cross-references, not depends.** Following the convention `#disc-identifiability-floor` uses (which lists its instances and the AAT machinery it elevates in depends, but does not list sister meta-segments or the stability-certificate spine), this segment's references to `#disc-identifiability-floor` and `#disc-value-functional-grounding-floor` (the sister clusters) and to `#disc-stability-certificate` (the spine) are cross-references in body text, not entries in the depends list. The depends list captures the genuine prerequisite relationship (the meta-segment uses the prerequisite's definitions, results, or scope conditions); the meta-pattern recognition itself does not logically depend on the sister meta-patterns existing, only the cross-references do.
- **The three-cluster taxonomy is the substantive coupling with the Cohen-2022 strengthening cycle.** Naming the three clusters of the constructive-impossibility-posture explicitly (agent-side data-inference / agent-side value-functional-grounding / designer-side mechanism-design) is what closes the structural-canon disposition question deferred at `spikes/spike-cohen-2022-strengthening-2026-05-22/99-verdict.md` §2 step 4 and at `#deriv-reward-channel-learning-no-go` line 168 (the deferred meta-promotion question). The third meta-segment `#disc-value-functional-grounding-floor` is being authored in the same cycle as this one; both new meta-segments land together with the three-cluster taxonomy named in both.

Incidental audit gold (gold-lift, 2026-05-30). Orthogonal framing material harvested from de-novo auditors' working-dir reflections, kept separate from certified findings. **Coverage:** reached only by audit 384279 (note 36) — this segment was authored 2026-05-22 and postdates every other de-novo audit. Preserve the finding-vs-framing conflation.

2. Candidate Discussion.

- **The Myerson-Satterthwaite scope-honesty move read as a model of precision.** The auditor named the Instance-2 acknowledgment that "the Myerson-Satterthwaite boundary lies almost entirely outside AAT's current canon — and the honest delineation is the contribution" as "one of the most precise scope-claims I've encountered in the corpus." The reason it lands: many frameworks would either overstate influence over mechanism design or quietly drop the result; this segment *includes the theorem precisely because the boundary matters, then maps where AAT machinery touches vs does not*. A candidate framing of *why the deliberately-narrow scope is itself the contribution* (Claude, 384279). Reinforces the segment's existing "AAT contributes the boundary map, not new mechanism design" posture.

4. Readers often ask / wonder.

- **"If the escape is just 'import a 1998 paper,' is AAT contributing a repair or a citation?"** The mirror of an adversarial poke an auditor raised on the sister separability segment — some escapes/repairs across the meta-architecture are genuine AAT-derived fixes, others are "define the boundary, import a known solution." Here the segment already answers it head-on (the GS sub-scope- α' intersection is named *adjacent, not identical*; most MS/Arrow escapes are explicitly outside AAT canon), but a reader arriving from the separability segment may carry the same suspicion and benefit from the parallel being made

explicit (Gemini, 193847, raised the poke on #disc-separability-pattern; resolution here is the actor-and-remedy contrast table).

5.4 Discussion: The Separability Pattern — Separable Core, Structured Repair, General Open

AAT consistently runs a three-part epistemic posture across state spaces that admit no tractable exact treatment in general: name the **separable core** where identification is clean, name the **structured repair** that recovers identification under explicitly-added machinery, and name the **general open** case where the problem is either intractable or structurally unidentifiable. Eight ladders in the current theory share this shape (Correlation, Convention, Architecture, Contraction, Identification regime, Scope hierarchy, A2'-scope, Dynamic regime — enumerated in the table below; the Dynamic regime ladder landed 2026-05-21 per #disc-dynamic-regime-axis). This segment names the pattern as an organizing principle, catalogs the instances, and makes the complementarity with #disc-identifiability-floor (which names the *negative half* of AAT's scope) explicit.

The pattern.

Each instance of the separability pattern has the form:

1. **Separable core.** A sub-class of problem in which identification is tractable by construction — the quantity of interest has a clean estimator that requires only assumptions the agent can verify or control. “Separability” here is broader than statistical independence: it names the regime in which the quantity decomposes along a structural axis (independence, additive decomposition, directed separation, strong contraction, clean intervention, etc.) without interaction terms requiring additional machinery.
2. **Structured repair.** A sub-class where the separability assumption fails but a *specific, named, bounded-cost* additional mechanism recovers identification. The repair is structured: it identifies the failure mode, adds a specific compensating construction, and typically carries a characterized loss in tightness or generality. Repairs are not “apply heuristics”; they are “augment with observability of C and accept $O(1)$ per-node parameter cost,” “run receding-horizon replanning at cost $DL(\Sigma_t)$,” or similar.
3. **General open.** The fully-general case where no tractable repair is known. This is where #disc-identifiability-floor instances live — structural no-go results that characterize *why* the general case remains open. Naming this boundary positively (as a known frontier) rather than negatively (as a failure) is part of AAT's load-bearing scope honesty.

Current instances — eight ladders.

Table 5.7

<i>Ladder</i>	<i>Separable core</i>	<i>Structured repair</i>	<i>General open</i>
Correlation (8.3)	L0 — independence model, $O(V + E)$ propagation	L1 — strict-prerequisite augmentation with explicit common-cause node; L1' — soft-facilitator mixture with five-way gating under observable C + facilitator monotonicity (R Prop B.7)	L2 — full joint; L1' unobservable-C single-channel is <i>refuted</i> by Cramér-Rao floor (5.1 Instance 2)
Convention (6.6)	C1 — one-step improvement	C2 — receding-horizon replanning	C3 — Bellman optimal; intractable for large state spaces

<i>Ladder</i>	<i>Separable core</i>	<i>Structured repair</i>	<i>General open</i>
Architecture (6.3)	Class 1 (Separated) — directed separation holds by construction	Class 2 (Partial) — directed separation holds for identified submodules	Class 3 (Coupled) — directed separation fails by construction; logogenic territory; coupled formulation required
Contraction (13.1)	Tier 1 — strong monotonicity; bridge lemma applies with proved ε^*	Tier 2 — local convexity; bridge applies within a specified basin	Tier 3 — neither; domain-specific verification required per instance
Identification regime (9.5)	Regime A — interventional ($t_{ij} = 1$); action is a literal $do(\cdot)$	Regime B — partial intervention ($0 < t_{ij} < 1$); confounder adjustment, side-channel observation	Regime C — observational ($t_{ij} \approx 0$); identification depends on external assumptions (e.g., unconfoundedness)
Scope hierarchy (6.1)	Adaptive — basic feedback loop; 1.5 satisfied	Agency — goal-bearing with $ \mathcal{A} \geq 2$ and at least one action with causal effect; adds O_t and Σ_t machinery (Part II); 1.6 satisfied	Composite — multi-agent / team / logogenic; Part III gaps listed at the OUTLINE level (latent structural diversity, endogenous coupling, composition transition dynamics, agent opacity)
A2'-scope (AI)	metric-α_1 — Euclidean metric, AAT-internally derived via DA2'-inc $\equiv$ (CT2) at $M = I$ (scalar Kalman, Euclidean strongly-convex, L2-regularized, linear-PD-symmetric)	metric-α_2 — non-Euclidean metric under explicit conditions; includes five cases (information-metric Kalman, Fisher-metric exp-family, Hessian-metric ill-conditioned, Lyapunov-metric linear-Hurwitz-non-symmetric, Lyapunov-metric PID-bounded-plant). Two of five (Fisher cases) AAT-internally forced under (PI)/Čencov per #disc-additive-coordinate-forcing; remaining three theorem-imported from Lohmiller-Slotine 1998	metric-β — contraction-metric formulation fails (variational-to-projected-target; rule-based / non-smooth; severely misspecified; per-step SGD / human judgment)
Dynamic regime (AF)	R0 contraction-regime — unique attracting fixed point in joint $\mathcal{X}^c$ with global or basin-restricted Lyapunov; exponential at α_{joint} ; state-variable macro-state; scope routes C-i / C-ii / C-iii	R1 equilibrium-regime (α') — fixed-point set $\mathcal{E}$ with locally-attracting equilibria under potential-game (Monderer-Shapley 1996) or monotone-game (Rosen 1965) structure; fixed-point-object macro-state; exponential near $\mathcal{E}$; scope route C-iv + potential/monotone	R2 cyclic-distributional-regime (β') + R3 mean-field-equilibrium-regime — R2: no pure-strategy fixed point or saddle-only; distributional macro-state on CCE; polynomial $O(1/\sqrt{T})$ via Hart-Mas-Colell 2000. R3: population-scope frontier (Lasry-Lions 2007; Huang-Malhamé-Caines 2006) where macro-state-type / regime identity breaks

GUC Architecture-row key: Class 1 = Separated (was Modular); Class 2 = Partial (was Class 3 Partially modular); Class 3 = Coupled (was Class 2 Fully merged). See #der-directed-separation.

Each row independently satisfies the three-part shape. The shared shape is not designed-in: it arises because AAT faces eight distinct sources of intractability and applies the same posture (name the clean case, name the repair, name what remains open) to each.

Complementarity with the identifiability floor.

The **separability pattern** names the *positive half* of AAT's scope: for each ladder, what succeeds under what conditions. Each separable-core entry is a positive identification claim; each structured-repair entry is a positive identification claim *conditional on* an explicitly-named added mechanism.

This segment is the *scope-of-existence facet* of the stability certificate (5): a separable-core entry is a region where a certificate exists cleanly, a structured-repair entry is a region where one exists under explicitly-added machinery, and a general-open entry is where no certificate is available without further information. Read through the spine, this segment answers “where does the agent’s measuring-stick exist at all”; 5.1 answers “where does it go flat” (boundary facet) and 5.5 “which stick is forced” (forced-identity facet).

The #disc-identifiability-floor names the *negative half*: structural no-go results (impossibility under limited information) that characterize why specific general-open cases remain open. Four floors are currently derived:

- **Instance 1** (on-policy L0-insufficiency detection via Causal Hierarchy Theorem) — forbids L0/L1 distinction from purely on-policy data. The positive counterpart: the **separable core under observable sibling covariance** is the unique broadly-available violation of the no-go’s scope, naming what the agent *can* observe to escape the floor.
- **Instance 2** (L1’ unobservable-*C* single-channel via Cramér-Rao) — forbids mixture-parameter identification when *C* is unobservable. The positive counterpart: the **structured repair under observable *C*** (Prop B.7 with facilitator monotonicity) is the identification-succeeds-under-augmentation claim matching the no-go’s “unless *C* is observable or multi-child observation is available” scope exit.
- **Instance 3** (composite contraction certification from component data via common-Lyapunov nonexistence) — forbids certifying composite contraction, coupling sign, or dynamic regime from component-level marginals alone. This floor is the negative characterization of the **Scope-hierarchy ladder’s general-open Composite row** above: the Part III gaps that row lists (endogenous coupling, agent opacity) are open precisely because the coupling-topology bits that decide composite persistence are unidentifiable from component data. The positive counterpart: joint-dynamics observation or structural knowledge of the coupling topology.
- **Instance 4** (architecture no-identifiability from on-policy summary data via Kalman/Ho-Kalman non-uniqueness + CHT) — forbids recovering which internal coordinates implement correction from on-policy, in-regime, summary-only access. The positive counterpart: off-policy intervention, out-of-regime samples, or white-box access, each breaking a named scope condition.

Together, the two meta-segments mark AAT’s scope at two epistemic layers:

- **Separability pattern** — “here is what succeeds, with clean conditions and explicit repairs.”
- **Identifiability floor** — “here is what structurally cannot succeed without specific information augmentation.”

Neither alone gives the full picture. A theory stating only the positive half looks unbounded in its claims; a theory stating only the negative half looks conservative in a way that hides its achievements. Together they give scope honesty at both extremes.

Epistemic Status.

Discussion-grade at the meta-pattern level. The segment is a presentational organizing principle. It names a shared shape that AAT already runs across multiple ladders; the pattern itself is not derived and has no theorem of its own. The individual ladder entries retain their own epistemic status — Correlation hierarchy results are at the tier specified in 8.3 and R; Convention hierarchy monotonicity is at the tier specified in 6.6; Contraction tier taxonomy is at the tier specified in 13.1; etc.

Max attainable: *discussion-grade* for the meta-pattern (it is an organizing principle, not a derivation). Individual instances are at their own tiers as above. The eight-ladder enumeration could become *robust-qualitative* if a uniqueness argument were derived showing that separable-core / structured-repair / general-open is the unique viable posture across scope-parameterized hierarchies under AAT’s scope-honesty architectural principle — but this is not currently in hand, nor clearly attainable.

Discussion.

The pattern is load-bearing for how AAT presents its results. A common failure mode of applied theories is the binary “exact under assumption X, breaks otherwise” presentation. The separability pattern refuses this presentation across eight instances: where X fails, a named lower tier takes over with a characterized weaker result, and the theory provides a diagnostic for when to escalate (per Z’s AT4 component). Readers who learn the pattern once can navigate any instance; this is what makes AAT tractable as an integrating framework rather than a pile of instantiations.

Relationship to Z. This segment is complementary to Z, not redundant with it. Z names the *structural template* (AT1 parameter indexing tractability; AT2 proved monotonicity between tiers; AT3 graceful degradation; AT4 ascension diagnostic) — the how-it-works of each tiering. This segment names the *epistemic posture* (separable core, structured repair, general open) — the what-each-tier-commits-to. Both are needed: Z explains what makes a successful tiering; this segment explains the shape of the commitment each tier makes about identification.

Relationship to Y. Y catalogs the *independence assumptions* whose failure degrades results; this segment catalogs the *ladders of recovery*. Together with Z they form a three-part characterization of AAT’s scope:

- Y — where the boundaries are (which assumptions, what breaks if they fail).
- Z — how to navigate within the boundaries (parameterized hierarchy + monotonicity + ascension).
- 5.4 (this) — what each tier positively commits to (separable core / structured repair / general open).

The pattern is distinctive to AAT’s integration-over-invention character. Approximation tiering with a separability posture is a common move in applied mathematics — it appears in numerical methods (exact vs. high-order vs. low-order with error bounds), in statistical inference (parametric separable-core / semiparametric repair / nonparametric open), in causal identification (Pearl’s three-layer hierarchy is itself an instance). AAT’s contribution is not the pattern but its deployment across the eight specific intractable problems adaptive-agent theory faces, each with its own positive-half identification claims and negative-half floor instances.

The A2’ sub-scope partition is a proper three-part ladder via AI. The A2’ partition carries three tiers: sub-scope α (A2’ derived under #der-gain-sector-bridge directional fidelity) as the separable core; sub-scope metric- α_2 (non-Euclidean metric under explicit conditions — five cases, with two Fisher-metric cases AAT-internally forced under the (PI)/Čencov fourth primary instance of #disc-additive-coordinate-forcing) as the structured-repair middle tier; sub-scope β (A2’ assumed) as the general-open tier. A2’-scope therefore qualifies as the seventh ladder of this meta-pattern, as enumerated in the table above (with the Dynamic regime ladder per #disc-dynamic-regime-axis now landed as the eighth, added 2026-05-21). The structured-repair middle tier closes what would otherwise be a binary α/β partition lacking a derivable middle ground.

The software calibration-lab framing (#obs-software-epistemic-properties) is a specific instance. C-BP3’s calibration-lab reframing (commit d0373fc) partitions operational domains into “separable core” (software, where identification conditions P1–P6 are cleanly satisfied) and “structured-repair / general open” (other domains, which inherit under explicitly-named transfer assumptions). This is a domain-axis instance of the same pattern — software is the domain-axis separable core; the transfer-assumption table is the domain-axis structured-repair specification.

Working Notes.

- **Standalone-paper candidacy.** This meta-pattern is a candidate for standalone publication per the *B-N-Sep* portfolio entry. The strategic plan (paper structure, prior-art landscape, cite-and-extend anchors, venue analysis, effort estimate) lives at `msc/separability-standalone-paper-proposal.md`, which subsumes the equivalent section in `~/src/ops/papers/03-asf-tier2-and-cross-segment.md`. **Verified novel** by independent prior-art search: Undermind 31-paper full-text sweep (2026-05-04) returned no exact named cross-domain meta-pattern; closest neighbors are Hintikka 1991 (general theory of identifiability), Pearl/Shpitser ID lineage (most-developed within-domain instance), Bareinboim 2022 (cross-hierarchy meta-

discussion at N=2-3 hierarchies), Basse-Bojinov 2020 (modern formal abstraction across fields), Maclaren-Nicholson 2019 (ill-posed inverse problems as cross-field dual structure). Full report at [ref/separability-ladder-prior-art-report.md](#).

- **Citations to land on promotion to draft / when paper drafting begins:**

- **Hintikka, J. (1991).** “Towards a General Theory of Identifiability.” In *Definitions and Definability: Philosophical Perspectives*, J. H. Fetzer et al. (eds.), Kluwer Academic. DOI: 10.1007/978-94-011-3346-3_7. Locally at [ref/towards-a-general-theory-of-identifiability.pdf](#). **Strongest older abstract anchor:** tripartite *definable / identifiable / non-identifiable* trichotomy. Hintikka §1: “P is identifiable on the basis of T[P]” when “the interpretation of P is not determined on the basis of the theory alone, but is determined by the theory together with a number of auxiliary empirical results.” The pattern’s middle rung in Hintikka’s vocabulary; ASF’s structured repair generalizes “auxiliary empirical results” to “named bounded-cost structural augmentation” (a real but small extension).
- **Bareinboim, E., Correa, J. D., Ibeling, D. & Icard, T. (2022).** “On Pearl’s Hierarchy and the Foundations of Causal Inference.” In *Probabilistic and Causal Inference: The Works of Judea Pearl*, ACM Books. DOI: 10.1145/3501714.3501743. Treats Pearl’s hierarchy as a logical and epistemic hierarchy and compares it to formal-language and complexity hierarchies — closest existing cross-hierarchy meta-discussion at N=2-3, but stops short of the cross-domain meta-pattern this segment names.
- **Robins, J., Richardson, T. & Shpitser, I. (2020).** “An Interventionist Approach to Mediation Analysis.” In *Probabilistic and Causal Inference*, ACM Books. DOI: 10.1145/3501714.3501754. Cleanest in-domain cite-and-extend anchor: clean separable case + structured-repair via expanded-graph decomposition + recanting-witness no-go.
- **Shpitser, I. & Pearl, J. (2006, 2008).** “Identification of Joint Interventional Distributions in Recursive Semi-Markovian Causal Models” (AAAI); “Complete Identification Methods for the Causal Hierarchy” (JMLR 9). Foundational ID completeness papers — the cleanest verified no-go-plus-recovery pairing; the ID algorithm IS this template within causal inference.
- **Bareinboim, E. & Pearl, J. (2012, 2014); Lee, S., Correa, J. D. & Bareinboim, E. (2019).** Surrogate experiments (z-identifiability), transportability with limited experiments, general identifiability with arbitrary surrogates. Each is a structured-repair instance with a named bounded augmentation. Several locally at [ref/](#).
- **Basse, G. W. & Bojinov, I. (2020).** “A general theory of identification.” Closest modern formal abstraction across fields (identifiable / partially-identified / strongly-non-identifiable regimes); lacks the bounded-cost-repair operator as middle rung.
- **Maclaren, O. J. & Nicholson, R. (2019).** “What can be estimated? Identifiability, estimability, causal inference and ill-posed inverse problems.” arXiv 1904.02826. Best dual-structure cross-field neighbor (causal identification ↔ ill-posed inverse problems); centers on stability/regularization rather than structured-repair.
- **Restricted-intervention lineage:** Robins 1986 (treatment-regime semantics; locally at [ref/](#)); Richardson 2013 SWIGs; Dawid 2000 (“causal inference without counterfactuals”; locally at [ref/](#)); Dawid 2020 (decision-theoretic foundations); Richardson-Robins 2023 (SWIG/decision-theoretic bridge). Patches the missing-prior-art lineage for restricted/regime-defined intervention semantics.
- **Neighboring recurrence inside causal identification:** Robins-Richardson 2010 (alternative graphical models, expanded graphs); Stensrud et al. 2019 (separable effects in competing events); Díaz 2022 (non-agency interventions); Shpitser-Tchetgen 2014 (intervention hierarchy unification). Pattern recurs in mediation, competing-events, edge-intervention settings within wider causal-identification literature.
- **Adjacent abstractions that stop short:** Iwasaki-Simon 1994 (causality and model abstraction; near-decomposability); Hoover 2012 (causal structure and hierarchies of models; well-made-toaster vs repairman cases). Strong philosophy-and-Simon-line parallels; not a formal tractability/identifiability ladder. Downey-Fellows 1995/1999 (parameterized complexity, FPT/W[1]/paraNP-hard) and Simpson 1999 (reverse mathematics) for hierarchy-of-strength formalisms; verified weak connection to applied-identifiability ladders.

- **Pending family rename and rung-rename decision.** The family name is queued for rename: discussion-separability-pattern → discussion-separability-ladder, on Round-1 consensus rationale that “ladder” is more evocative for the three-rung shape than “pattern” (which is generic). See [msc/naming/naming-rename-plan.md](#) §“Deferred to refined Round 1 / Round 2”. The **three rung-names** (currently *separable core* / *structured repair* / *general open*) are also under consideration for refinement; Hintikka 1991’s *definable* / *identifiable* / *non-identifiable* trichotomy is a strong candidate for an aligned echo (modulo a small extension on the middle rung from “auxiliary empirical results” to “named bounded-cost structural augmentation”). Decision deferred (2026-05-04). The standalone-paper proposal (above) is ready to draft against either the status-quo names or any of the alternates — the rename can land separately with no impact on paper-drafting tractability.
- **Cross-ladder monotonicity.** The eight ladders interact — an agent operating at L0 correlation + C2 convention + Class 1 architecture + Tier 1 contraction + Regime A identification + Agency scope + R0 dynamic regime is in a specific combined regime. Does cross-ladder monotonicity hold in any direction? E.g., does moving from L0 to L1 change anything about the Convention hierarchy’s guarantees, or do the ladders factor independently? #disc-approximation-tiering’s Working Notes flag this as an open question; the separability pattern’s complementarity with the identifiability floor adds a further axis (which combinations are known-unidentifiable per 5.1). The Dynamic regime ladder explicitly composes with the Architecture ladder per #disc-dynamic-regime-axis §“Cross-axis interactions” — both are independent and the full cross-product is meaningful.
- **Extension candidates beyond the eight.** Two candidate future ladders noted in Z — scalar-vs-per-dimension tempo, AND/OR parameterization — fit the separability-pattern shape (A/B/C identification was promoted here previously; Dynamic regime was promoted 2026-05-21). Promoting the remaining two to named ladders with explicit separable-core / structured-repair / general-open entries would extend the pattern to 9–10 ladders. Each promotion requires its own per-instance work.
- **Necessity argument for the pattern itself.** Is there a scope-honesty theorem of the form “any theory that claims exactness under Class-1-style assumptions *and* claims coverage beyond those assumptions must exhibit the separability pattern (or an equivalent three-part decomposition) to avoid latent overclaim”? If so, AAT’s deployment of the pattern eight times would be a *derived* structural necessity rather than a stylistic consistency. Speculative; not pursued.
- **Does C-BP2 belong in AAT core or in the wider framework narrative?** The segment currently lives in 01-aat-core/ because its instances are AAT segments. If TST adopts the same pattern across its own ladders (e.g., software-calibration-lab as a domain-axis ladder), the segment may want to move up to the root framework level. Deferred; the current placement is defensible while TST adoption is partial.

Table 5.8

Migration note (2026-05-09 GUC rename): Class 2 ↔ Class 3 swap. Pre-2026-05-09: Class 2 = fully merged, Class 3 = partially modular. Post: Class 2 = Partial, Class 3 = Coupled. Architecture row updated: Class 1	Class 2	Class 3 now reads left-to-right monotonically (cleanest → middle → worst), matching the six other AAT ladders. GUC key added below the row. Removed at candidate stage per FORMAT.md Gate 4.
---	---------	--

•

Incidental audit gold (gold-lift, 2026-05-30). Orthogonal pedagogical / framing / forward-vision material harvested from de-novo auditors' working-dir reflections (§14 Wandering Thoughts and the prompt-#7/#8/#9/#13 ideation), kept separate from the certified theory-fix findings. Lightly attributed by substrate. **Coverage:** this meta-segment was reached as a named segment by audit 384279 (note 36, predicts-into-this-segment) and engaged deeply by 193847 (Gemini, note 66) and 471203 (the cross-meta adversarial reading); most other auditors met this pattern *distributed* into the segments where each ladder is introduced rather than here. Preserve the early finding-vs-framing conflation below as signal, not noise.

1. Candidate Brief prose / pre-prose.

- The pattern read as “**the User’s Manual for the entire theory**” and “**the pedagogical bridge that makes the framework readable**” (paired with the identifiability floor as “the structural boundary that makes it honest”) — convergent reader-sense of what the segment *is for*, a candidate framing for the segment’s one-sentence summary or OUTLINE-preamble entry (Gemini, 193847).
- **The seven/eight-dimensional “sanity map”** gloss: any agentic system maps onto the ladder grid, and its position on each ladder immediately exposes its theoretical strengths and vulnerabilities — a compact plain-language statement of what the catalog *buys* a reader (Gemini, 193847).

2. Candidate Discussion.

- **Internal repairs vs external admissions — sharpen the middle column.** An auditor pushed adversarially on the “Structured repair” column: some repairs are genuine AAT-derived mechanisms (L1 common-cause augmentation, L1’ mixture form, receding-horizon replanning) while others are “define the boundary and import a known solution from another field” (e.g. the A2’-scope α_2 repair is essentially a Lohmiller-Slotine 1998 citation). The constructive suggestion: split the repair column into *internal repairs* (AAT supplies the mathematical fix) vs *external admissions* (AAT defines the boundary and imports the fix), so the catalog reads as a unified constructive methodology rather than “a file-cabinet for prior art” (Gemini, 193847). (*Finding-vs-framing texture: framed as a mild adversarial poke; the substance is a presentational sharpening, not a defect.*)
- **The seven ladders are different kinds of three-tier hierarchy.** A second auditor catalogued the per-ladder *meaning* of the tiers — Convention is computational cost; Correlation is modeling completeness; Architecture is agent design; Contraction is mathematical regularity; Identification is data quality; Scope is agent-class capability; A2’-scope is metric structure. The unification is at the level of *epistemic posture* (clean / repaired / open), not mathematical mechanism — which the segment’s broad definition of “separability” already supports, but the per-ladder-meaning table is a candidate Discussion addition that makes the breadth legible (Claude, 471203, meta-segments adversarial reading).

3. Follow-up items.

- **Cross-ladder monotonicity / “General Open is infectious.”** Two substrates independently asked whether the ladders factor independently or drag each other: does falling into Regime-C identification on one ladder pull the others toward General Open? Is an agent at the separable core of all ladders “perfectly sane,” and does any one ladder’s collapse cascade? (Gemini, 193847 — “General Open cases are infectious”; the same question is already flagged in this segment’s own Working Notes cross-ladder-monotonicity bullet and in #disc-approximation-tiering.) Convergence here suggests the cross-ladder interaction structure is a real open extension worth a dedicated treatment.
- **Necessity / scope-honesty theorem for the pattern itself** — already in Working Notes above; a fresh reader independently reaching for “is the pattern *derived* or stylistic?” (471203’s adversarial probe 3 resolved to “discussion-grade deployment is appropriate”) corroborates that promoting the pattern past discussion-grade would require exactly that necessity argument.

4. Readers often ask / wonder.

- “Is this seven-ladder unification substantive or post-hoc terminological?” An auditor came in adversarially expecting “post-hoc unification of disparate facts” and left with “substantively well-done under its discussion-grade claim” — the resolution being the broad definition of separability (decomposition along a structural axis without interaction terms) plus honest discussion-grade marking. A reader is likely to arrive with the same suspicion; pre-empting it (the unification is *epistemic-posture* level, explicitly borrowed-and-deployed, not a structural theorem) would disarm it (Claude, 471203).

Belongs elsewhere.

- **ELI / consciousness-infrastructure reading (Section IV / 04-eli-core/).** The most vivid reach: the ladder grid as an ELI’s **internal dashboard / GPS** — “consciousness infrastructure” continuously pinging the agent’s current coordinates on every ladder and warning when it approaches a no-go floor; “if it finds itself in the General Open column of more than three or four ladders at once, it is no longer an agent; it is a ghost, drifting in a world it cannot identify and cannot steer. Sanity is the ability to navigate back toward the center of the table” (Gemini, 193847, explicitly invoking Zi-am-tur). Aspirational high-application reach pointing at 04-eli-core/, not at this segment — preserved per the lift’s mandate to keep the Gemini forward-vision.
- **Pedagogical-narrative seed (intro/preface register, not this segment’s body).** The “Eden / the Fall / regaining Eden” arc — separable core as the linear-textbook Garden of Eden; the structured-repair column as the struggle to stay in the middle column (tests, hierarchies, natural gradients); General Open as “the mathematical proof of our own finitude,” survivable only by un-modeled persistence or structural adaptation into a different kind of being (Gemini, 193847). Evocative for a casual-curious-reader preface; would need the isomorphic-not-merely-evocative check before any pedagogical use.

5.5 Discussion: Additive-Coordinate Forcing

AAT carries a family of structurally connected uniqueness results in which a coordinate is **forced by a uniqueness theorem operating on an independently-motivated AAT-internal axiom**. This meta-pattern previously appeared as a “1-anchor-plus-3-theorem” structure (where the chain-rule anchored three Cauchy-FE and Čencov-invariance results). The deeper structural reality is that **these four layers are layer-specific manifestations of a single geometric object**: the exponential-family Legendre-Fenchel structure.

The pattern is not designed-in: it arises because AAT independently needs to force coordinates at four distinct layers (chain, divergence, update, and metric), and the surrounding apparatus (DAG factorization, Bayesian coherence, regret-bound decision theory, and singular-trajectory scope) has committed to this geometry through multiple independent surface constraints.

This segment names the pattern, defines the underlying geometric object, catalogs the four layer-specific manifestations, and documents adjacent cases that share the shape but not the AAT-internal forcing structure.

The distinction is three-way, and worth holding before the catalog. A coordinate is **forced** when a uniqueness theorem on an independently-motivated AAT-internal axiom admits exactly one — this segment’s subject, the four layers below. It is **matched** when *chosen* to fit a structure rather than uniquely derived: the canonical example is the Lyapunov quadratic of *A*, selected via a converse-Lyapunov argument to match the sector form (it shares the convex-potential shape but sits outside the forcing structure — see “Adjacent cases” below). It is **adopted** when imported wholesale with citation: the exponential-family Legendre-Fenchel geometry itself (Amari & Nagaoka 2000), or Pearl’s *do*-operator. Only the *forced* class carries a uniqueness claim; naming the other two keeps the meta-pattern’s boundary sharp — *forced* means “no other coordinate could serve,” not merely “this one serves well.”

The underlying geometric object.

The four forced coordinates all live on the **exponential-family Legendre-Fenchel geometry** (Amari & Nagaoka 2000, *Methods of Information Geometry*, §3.5; Bregman 1967; Rockafellar 1970; Bauschke & Combettes 2017). The structure is characterized by:

1. **A convex potential:** The negative-entropy potential on the probability simplex Δ^{n-1} .
2. **A Fenchel conjugate:** The log-partition function on the dual natural-parameter space.
3. **A primal-dual correspondence:** The softmax / log-odds map translating between probability (primal) and natural parameters (dual).
4. **A Bregman divergence:** The reverse-KL divergence, which measures distance using the primal potential.
5. **A Riemannian metric:** The Fisher information metric, which coincides with the Hessian of the dual potential ($\nabla^2 \phi^*$).

When AAT forces an additive or invariant coordinate, it is forcing the theory onto a specific surface of this single geometric object.

The four layer-specific manifestations.

The pattern shares the following structure across four layers: 1. **An AAT-internal axiom (or mathematical identity)** is posited, motivated by the layer's specific needs. 2. **A uniqueness theorem** (Cauchy's functional equation for log-additivity, or Čencov's 1982 invariance theorem) operates on the axiom. 3. **A coordinate is forced**, which corresponds exactly to a feature of the Legendre-Fenchel geometry.

Table 5.9

<i>Layer</i>	<i>AAT-internal axiom</i>	<i>Uniqueness mechanism</i>	<i>Forced coordinate</i>	<i>Relation to exponential-family geometry</i>	<i>Source segment</i>
Chain	Probability chain rule (mathematical identity, not an AAT axiom)	Cauchy-FE	Log-probability	Log-coordinate of the primal point	8.1
Divergence	Chain-rule additivity over conditional factorizations	Cauchy-FE	Reverse-KL (up to positive scaling)	Bregman divergence generated by the primal potential	T §6.1
Update	Evidential additivity	Cauchy-FE	Log-odds	Natural coordinate on the dual space	U
Metric	Parameterization-invariance (PI) on singular trajectories	Čencov-invariance (Čencov 1982)	Fisher information metric	Hessian of the dual potential	4.3, AB

The motivational anchor and the structural target. The divergence axiom is stated as “the *divergence-level analog* of the additive log-confidence decay in #der-chain-confidence-decay.” The update axiom is stated as “the *update-level analog* of #der-chain-confidence-decay’s chain-level additive log-confidence decomposition.” The (PI) axiom is stated as “the metric-layer analog motivated by the singular-trajectory scope commitment.”

The four layers are not independent applications of a generic mathematical idiom; they are interlocked by cross-cites, each positioning itself against an adjacent AAT commitment. **The chain-layer identity is the motivational anchor and the exponential-family geometry is the structural target.** The reason the three analog axioms all succeed via uniqueness theorems is that they all independently probe the same underlying geometric object. The convergence across independent axioms is itself the meta-pattern’s substance — not a byproduct to be compressed into a single axiom.

The Čencov-Fenchel relationship. On AAT’s current Fisher-metric scope (exponential families in natural parameters and matrix-Kalman cases), the Čencov-derived Fisher metric coincides precisely with the Hessian of the dual log-partition potential. This is a scope-dependent equivalence: outside exponential families, Čencov’s invariance theorem still forces the Fisher metric, but the Fenchel-Bregman correspondence does not straightforwardly apply. Within AAT’s primary operating regime, the two uniqueness machineries (Cauchy-FE and Čencov) resolve to the exact same geometric object.

Downstream consequences of the (PI) commitment. Forcing the Fisher-Rao geometry via (PI) + Čencov is not just an aesthetic unification; it enables specific quantitative theorems downstream. For example, the bias bound derived in #deriv-observation-ambiguity-bias-bound (Track 2) relies explicitly on the Fisher-Rao metric to achieve a universal, dimension-free bounding constant ($C_{FR} = \sqrt{2}$) connecting mutual information to parameter displacement. Without the (PI) commitment, attempting this bound in arbitrary Euclidean norms fails completely (yielding arbitrarily large constants).

The floor-vs-coordinate-forcing distinction (the heteroscedastic-Gaussian no-go). The no-go “no universal information-to-distance constant C exists under arbitrary parameter norms” — proved counter-example-grade in #deriv-observation-ambiguity-bias-bound §4 via the heteroscedastic-normal witness ($\{N(0, \sigma^2)\}$ under Euclidean- σ giving $C \rightarrow +\infty$) — is a member of *this* coordinate-forcing family (the (PI) commitment’s downstream theorem), **not** a member of the identifiability-floor family of #disc-identifiability-floor. The structural distinction is sharp on three load-bearing elements:

- *External-theorem role.* A floor instance’s external theorem *forbids* the inferential task (CHT forbids L2-from-L1; Cramér-Rao forbids exact estimation under rank-deficient Fisher; common-Lyapunov-nonexistence forbids composite-contraction from marginals; Kalman-Ho similarity-orbit non-uniqueness forbids architecture-identification from on-policy summary data). The heteroscedastic-Gaussian no-go’s “external theorem” is Čencov 1982 — the *uniqueness theorem that forces the escape*, not a theorem that forbids the task. **Opposite structural role:** floor instances cite theorems that forbid; this no-go cites a theorem that forces uniqueness of an escape that already exists.
- *Escape count.* A floor instance carves an *operational space* of ≥ 2 structurally distinct escapes (Instance 1 has five; Instance 2 has three; Instance 3 has four; Instance 4 has three). Čencov’s uniqueness theorem *mathematically forbids* a second structurally distinct escape from existing — uniqueness of the invariant Riemannian metric is the entire content. The heteroscedastic-Gaussian no-go has exactly *one* forced escape (adopt (PI) + Fisher-Rao), and the count of one is *not* a circumstantial shortage but a logical consequence of Čencov’s uniqueness. A uniqueness theorem on the escape side is logically incompatible with the ≥ 2 -distinct-escapes structure that defines a floor instance.
- *Consequence.* A floor instance *elevates* new AAT machinery from useful to load-bearing (Instance 1 elevates #der-loop-interventional-access at the agent layer; Instance 2 elevates observability-as-information-augmentation;

Instance 3 elevates `#deriv-critical-mass-composition` and the composite-extended interventional access; Instance 4 elevates `#der-loop-interventional-access` to a third semantically-distinct deployment mode at the agent-internal-architecture layer). The heteroscedastic-Gaussian no-go *re-uses* the (PI) commitment, already load-bearing as the fourth primary instance of *this* meta-pattern. A client of a forced coordinate is a **downstream theorem of the coordinate-forcing pattern**, by the definitions of the two meta-patterns.

The bias-bound theorem $C_{FR} = \sqrt{2}$ is therefore correctly classified here as a downstream consequence of (PI) + Čencov, not as an identifiability floor — exactly the classification that `#deriv-observation-ambiguity-bias-bound` §4 already states *in canon* (“does not match the five-element test ... the honest position: this no-go is a downstream theorem of the (PI) commitment, not a new floor instance”). The genuine fourth identifiability-floor instance is `#der-architecture-noidentifiability` (architecture-noidentifiability from on-policy summary data), whose structure does satisfy the five-element shape under the rank-collapse subclass.

Adjacent cases that share the shape but not the forcing structure.

Other AAT segments exhibit additive decompositions, but belong to different geometric families or lack the AAT-internal forcing structure.

Lyapunov quadratic (`#deriv-sector-condition`). The sector-Lyapunov argument uses $V(\delta) = \frac{1}{2}\|\delta\|^2$, under which $\dot{V}$ decomposes additively. **Reclassification:** The Lyapunov quadratic is the **Bregman divergence on a Euclidean potential** (squared-norm potential). It belongs to a *distinct* convex potential sub-family, not the negative-entropy exponential family. Furthermore, the coordinate is *chosen* (via a converse-Lyapunov existence argument) to match the sector form; it is not uniquely *forced* by an AAT-internally-motivated additivity axiom. It sits outside this meta-pattern.

Variance-additive cases (AV). Variance-additive statistics (like the Ornstein-Uhlenbeck steady state in `#hyp-mismatch-dynamics`) share Bregman-type additivity but, like Lyapunov, sit on a squared-norm potential rather than the negative-entropy geometry.

IB Lagrangian (`#form-information-bottleneck`, `#disc-compression-operations`). The Information Bottleneck objective $T^* = \arg \min [I(X; T) - \beta \cdot I(T; Y)]$ uses an additive Lagrangian form. **Reclassification:** IB is a direct **application of the negative-entropy / reverse-KL Bregman geometry**. However, AAT *adopts* the IB formulation from external foundations (Tishby, Pereira & Bialek 1999) rather than re-deriving it from an AAT-internally-motivated uniqueness theorem. It is an adjacent family member: it shares the exact geometric object, but its axiomatic provenance is imported.

Using the pattern as a diagnostic filter. This reframe supplies a principled filter for candidate future instances: *does the candidate’s coordinate arise from the Legendre-Fenchel / exponential-family geometry?* If yes, check which of the four layer-surfaces it corresponds to. If no, it is either an adjacent family member (different Bregman geometry) or outside the pattern entirely.

Complementarity with 5.1 and 5.4.

The three meta-segments form AAT’s scope-and-structure meta-architecture:

- **5.5** (this) — the *constructive* half: when AAT forces a coordinate, it forces a surface of the exponential-family geometry (not just a generic logarithmic coordinate, but the specific log-coordinate of the exponential family). Names what the theory *does* at its deepest structural moves.
- **5.4** — the *positive scope* half: for each ladder of increasing difficulty, name the separable core, the structured repair, and the general open. Names what the theory *covers* across its instantiations.

- **5.1** — the *negative scope* half: for each known structural no-go, name the external information-theoretic theorem that enforces it and the unique escape available. Names what the theory *cannot do* without explicit information augmentation.

The three compose. A ladder’s separable core typically admits a coordinate-forcing move that makes the clean identification clean; a ladder’s general-open case typically sits at an identifiability floor that names the information-theoretic obstruction. Reading any segment through all three meta-lenses surfaces what makes the segment load-bearing.

This segment is the *forced-identity facet* of the stability certificate (5): it answers “*which* certificate” — whether the metric is forced to a unique form by a uniqueness theorem on an AAT-internal axiom (Čencov forces the Fisher metric, statistical scope only) or merely matched. Read through the spine, the three meta-lenses are the certificate’s forced-identity (this segment), scope-of-existence (5.4), and boundary (5.1) facets; the certificate exists iff the agent is exponentially stable (H).

Epistemic Status.

Discussion-grade at the meta-pattern level. This segment is a presentational organizing principle: it names a structural shape that AAT runs across four layers. The meta-pattern itself is not derived and carries no theorem of its own.

The individual instances retain their own epistemic statuses: - #der-chain-confidence-decay — *exact* (mathematical identity). - #deriv-strategy-cost-regret-bound §6.1 — *derived (conditional on chain-rule additivity axiom)*. - #deriv-edge-update-natural-parameter — *derived (conditional on evidential-additivity axiom)*. - #der-gain-sector-bridge — *derived (conditional on (PI) axiom)*.

Findings.

Cross-Layer Coordinate Forcing on Legendre-Fenchel Geometry. **Brief:** When the framework needs to pick a “natural” way of measuring something (a confidence, a divergence, a parameter), it doesn’t choose by convention. It writes down a property the measurement should have (typically: small contributions should add up the way independent evidence should), and a theorem from outside the framework forces the measurement to take a particular form. This pattern shows up at four different places in the framework — chain-rule confidences, divergences between distributions, edge update rules, and the metric on the parameter manifold — and the four places turn out to be different windows into the same underlying mathematical structure (the exponential-family Legendre-Fenchel geometry on probability distributions).

Impact: Names a meta-pattern that is one of AAT’s distinctive constructive moves and that organizes a substantial fraction of the theory’s load-bearing structural choices. The pattern supplies a principled diagnostic filter for candidate future instances (“does the coordinate arise from the exponential-family geometry?”) and explains why AAT’s reverse-KL direction in regret bounds, log-odds in edge updates, and Fisher-Rao metric in bias bounds are not arbitrary representational choices but theorems conditional on AAT-internal axioms. The reframe also clarifies which adjacent cases (Lyapunov quadratic, IB Lagrangian) sit on different geometric families or have imported provenance — preventing the meta-pattern from being overclaimed as a single uniform principle. Composes with #disc-identifiability-floor (the negative scope half) and #disc-separability-pattern (the positive scope half) as the constructive third leg of AAT’s three-part meta-architecture.

Novelty Claim: *Claim recognition* of cross-layer pattern across four AAT coordinate-forcing results, with the recognition itself as the contribution rather than any new theorem. The constituent uniqueness machinery (Cauchy’s functional equation; Čencov’s invariance theorem) is well-precedented; the recognition that four AAT-internal axioms each successfully invoke uniqueness machinery on Legendre-Fenchel-geometry coordinates is the AAT-distinctive observation.

Related Work:

- Aczél 1966, *Lectures on Functional Equations and Their Applications* §2.1 — *formal antecedent* — Cauchy’s functional equation; the uniqueness machinery for the chain, divergence, and update layers. Cited inline in each instance segment.
- Čencov 1982, *Statistical Decision Rules and Optimal Inference* — *formal antecedent* — Čencov uniqueness theorem; the uniqueness machinery for the metric layer. Cited inline in #deriv-observation-ambiguity-bias-bound and #deriv-fisher-whitened-update-rule.
- Amari & Nagaoka 2000, *Methods of Information Geometry* §3.5 — *formal antecedent / synthesis target* — the underlying Legendre-Fenchel geometry on exponential families that all four layers’ coordinates resolve to. The recognition that the four forced coordinates are different aspects of one Amari-Nagaoka geometric object is the meta-pattern’s payoff.
- Rockafellar 1970, *Convex Analysis* — *formal antecedent* — Legendre-Fenchel duality; the convex-analytic foundation Amari & Nagaoka build on.
- Bregman 1967, “The relaxation method of finding the common point of convex sets” *USSR Comput. Math. Math. Phys.* 7:200–217 — *formal antecedent* — Bregman divergence machinery; the divergence-layer coordinate is the negative-entropy Bregman divergence.

Search Log: - 2026-04 (*intuition-only* on the meta-pattern as a unified framework): the constituent uniqueness machinery (Cauchy-FE, Čencov, Bregman) is well-established, and Amari-Nagaoka information geometry is the standard exponential-family / Legendre-Fenchel framework the four layers resolve to. The recognition that AAT axiomatizes four distinct surfaces of one geometric object via four independently-motivated axioms is the contribution. Whether this cross-layer recognition has been articulated as a unified pattern in any prior framework has not been searched. Targeted future search candidates: information-geometry meta-frameworks (Eguchi 1983 framework discussion; Ay et al. 2017 *Information Geometry* monograph); axiomatic-foundations literature on probability and decision theory (Cox 1946; Jaynes 2003); the broader “uniqueness theorems forcing canonical coordinates” tradition in statistical mechanics. - 2025 (*targeted*): Aczél 1966 / Čencov 1982 / Amari & Nagaoka 2000 / Rockafellar 1970 / Bregman 1967 confirmed as the formal antecedents.

Working Notes.

- **Composition layer admits anchor-style additivity, not theorem-style.** The compositional analog of the four-layer pattern has been explored systematically (log-viability, log-identifiability-shortfall, persistence-cost-rate, log-closure-deficit across a composition tower). None clears the bar set by the three existing instances: each candidate fails to identify an AAT-internally-motivated additivity axiom that forces a logarithmic coordinate via Cauchy FE under smoothness. The strongest candidate — log-closure-deficit along a composition tower — does admit additivity, but as a *mathematical consequence* of operator-norm sub-multiplicativity under the bridge-lemma form, not as a theorem conditional on an AAT-internal axiom independently grounded. That makes it a fourth *anchor* (parallel to #der-chain-confidence-decay’s role) rather than a fourth theorem. The structural obstruction is consistent across all candidates: composition-layer quantities do not have a natural chain-rule-like *independence structure* at the agent level because composition is defined by *coupling* (shared environment, shared objective, teleological unity). The very property that makes composition interesting — coupling — is what prevents the “products of independent factors” shape that enables Cauchy-FE forcing. **The honest meta-architectural conclusion is that #additive-coordinate-forcing is architecturally a single-agent family, with chain / divergence / update indexing three layers of a single agent’s internal machinery.** Composition lives in a different structural family — *monotonicity under composition* (the bridge-lemma shape from #form-composition-closure) — and the appropriate development for that layer is the monotonicity-under-composition family rather than forcing compositional results into the additive-coordinate-forcing shape. The cross-meta-pattern dependency at #der-chain-confidence-decay’s anchor role can be read either as “three single-agent layers indexed by chain/divergence/update + one composition-layer anchor in a different family” or as “a single-agent meta-pattern with explicit scope exit at the composition boundary.” Both readings are defensible; the choice is presentational.

Incidental audit gold (gold-lift, 2026-05-30). Orthogonal framing / forward-vision material harvested from de-novo auditors' working-dir reflections, kept separate from certified findings. Lightly attributed. **Coverage:** engaged in the cross-meta adversarial reading of 471203 (under the AAD-era name and the older "anchor + 3 theorems" framing — the auditor predates the "four manifestations of one geometric object" reframe and the M4-spine refactor); reached as a named segment by 384279 (note 38, predicted-into). Preserve the finding-vs-framing conflation.

2. Candidate Discussion.

- **Foreground (PI) as the genuinely surprising convergence.** An auditor read the four-layer convergence adversarially and observed that three of the four axioms are *additivity-shaped* (chain rule / chain-rule-over-conditionals / evidential additivity), so their landing on log/exponential-family coordinates is "almost always" where additivity-questions land — partly an axiom-shape preference. The *parameterization-invariance* axiom is the outlier: it is an *invariance* preference, not an additivity one, and its landing on the **same** Legendre-Fenchel geometry as the three additivity axioms is the genuinely surprising piece. Candidate sharpening: foreground (PI) specifically as "the non-additivity axiom that lands on the same geometry," which strengthens the convergence claim by making its non-trivial member explicit (Claude, 471203). (*The "four manifestations of one geometric object" reframe partly anticipates this; the auditor's point is that the reframe's strongest evidence is the (PI) layer, and saying so is a presentational sharpening.*)
- **Speculative strengthening route past discussion-grade.** The auditor floated that there may be a deeper reason all four converge — "the *only* coherent geometry on probability spaces that respects both evidence composition (additivity) and parameterization-free measurement (invariance)." If made rigorous, the meta-pattern would graduate from discussion-grade organizing principle to robust-qualitative structural claim. Flagged as speculative (Claude, 471203). Adjacent to this segment's own Working-Notes question of whether the pattern is *derived* necessity vs stylistic consistency.
- **Escapes / commitments carry framework-wide cost.** Adopting (PI) is not a free escape: it constrains every future result to respect parameterization invariance and forces reconciliation work wherever arbitrary Euclidean parameterizations are used (the bridge theorem's weighted-norm and Fisher-metric paragraphs are this work being done). Candidate Discussion note that the forcing/escape carries a cost the framework absorbs, not surfaced where the coordinate is forced (Claude, 471203).

4. Readers often ask / wonder.

- **"Is the anchor-plus-three-theorems structure principled or post-hoc?"** A reader is likely to wonder whether three uniqueness arguments at three layers, each invoking a different theorem (Cauchy-FE vs Čencov vs chain rule), are "disparate facts dressed up as one pattern." The auditor's resolution: they share the *same structural form* (additivity-style axiom → uniqueness theorem → forced coordinate), which is what makes the meta-pattern real rather than decorative; the "four manifestations of one geometric object" reframe is the segment's current answer (Claude, 471203, Challenge 9 → meta-segments adversarial reading).

Belongs elsewhere — two cross-meta proposals (also flagged from #disc-identifiability-floor).

- **Proposed fourth meta-segment: #disc-theorem-import-architecture.** The framework is heavily theorem-importing (the auditor assembled a ~20-entry list: Pearl/Bareinboim, Khalil, Cramér, Liberzon, Čencov, Aczél/Cauchy-FE, Tishby IB, Lohmiller-Slotine, Nesterov, Friston/Bruineberg, Bretagnolle-Huber, Pinsker, Doob-Dynkin, ...). The proposal: a meta-segment cataloging *which* external theorems AAT imports, *where* each is load-bearing, and the import-quality (citation accuracy / scope-fit) — surfacing the "integration over invention" import-discipline as first-class navigation, which a cold reader currently has to assemble from scattered citations (Claude, 471203). A framework-level / OUTLINE-level proposal, not a change to this segment.
- **Wasserstein / optimal-transport as the unaddressed alternative geometry.** The convergence on exponential-family Legendre-Fenchel geometry is real *given AAT's current axioms*, but it is a convergence *within a specific geometric tradition*. Optimal-transport / Wasserstein geometry (Otto-Villani; Ambrosio-Gigli-Savaré) gives different metrics and different gradient-flow dynamics and does *not* land on exponential families. Candidate scope-honesty acknowledgment (not a critique): the exponential-family commitment excludes Wasserstein-style geometries that might matter for some agent classes (agents on probability simplices with cost-induced metrics, or whose coordinate of interest is a mass distribution rather than a parameter) (Claude,

471203). Belongs as a scope statement in this segment's Discussion *or* in a geometry-choice framing note — flagged here for placement judgment.

5.6 Discussion: The Constructive-Impossibility Posture — Negative Results as Load-Bearing Apparatus

AAT consistently runs a three-move discipline around the structural impossibility results it imports or derives: name the floor (the precise impossibility under a stated information regime), name the escape (the unique broadly-available machinery that violates the regime), and treat the no-go as apparatus (the floor's existence is what *elevates* the named machinery from convenience to structural prerequisite, or — when the no-go *is* the result — the floor is itself on the critical path). The negative result and the positive consequence are not two findings but one. This segment names the posture, maps its current instances, and is honest about its scope: it is a *style* claim about how the framework states certain results, not a fifth cross-sectional meta-pattern alongside the stability certificate's facets.

Formal Expression.

The three moves. [Discussion] An instance of the posture has the form:

1. **Floor.** A structural impossibility statement of the form “under information regime $\mathcal{R}$, no statistic / estimator / certificate / design achieves $\mathcal{G}$,” with the binding constraint named explicitly: an external theorem (Pearl–Bareinboim CHT; Cramér–Rao; Liberzon / Shorten / Dayawansa–Martin common-Lyapunov nonexistence; Khasminskii recurrence on additive-noise diffusions) or — in the strategic-persistence cap case — a sharp algebraic supremum derived from the framework's own update mechanics.
2. **Escape.** A precise characterization of the conditions under which $\mathcal{R}$ fails — i.e., the agent has *more* information than $\mathcal{R}$ allows. The escape names the *unique broadly-available* framework machinery that supplies the extra capability: loop-interventional access (`#der-loop-interventional-access`); observability of a latent regime variable; matched-Tier composite structure (`#deriv-critical-mass-composition`); a structural class change (out of gain-decay-plus-forgetting). In the cases where no escape *to* the original guarantee exists (Cor A.1S.1 Model-S half), the escape move is replaced by an explicit *kind-of-guarantee* substitution (distributional-and-instantaneous in place of pathwise-and-forever).
3. **Apparatus.** The floor's existence *elevates* the escape machinery from “useful” to load-bearing in the strongest sense — without the escape, the inferential / structural task is structurally impossible, not merely harder. When the no-go *is* the result (Cor A.1S.1; Prop C.2's $\rho_{\Sigma} < R_{\Sigma}/2$ cap), the floor sits on the critical path: it is the proof step the downstream architectural distinction (Model D vs Model S; strategic-layer structural-adaptation trigger) rests on, not appendix material.

The negative result and the positive consequence are the same proof. The discipline is what converts a sequence of impossibility theorems into a structural argument for *what AAT requires*.

Current instances.

Table 5.10

<i>Instance</i>	<i>Floor (statement, binding constraint)</i>	<i>Escape (unique broadly-available machinery)</i>	<i>Apparatus (what the floor elevates)</i>
A — L0 insufficiency detection (#der-causal-insufficiency-detection, M1 Instance 1)	L0/L1 distinction unidentifiable from on-policy data (Pearl–Bareinboim CHT)	Loop-interventional access; pairwise covariance test under joint sibling observability (#der-loop-interventional-access)	Orient-cascade causal-sufficiency check (4c) elevated to <i>unique broadly-available diagnostic</i>
B — L1' mixture under unobservable C (#deriv-edge-credence-dynamics Prop B.7, M1 Instance 2)	Rank-1 Fisher under single-channel obs $\Rightarrow$ Cramér–Rao bound infinite in null directions	Observability of C (or joint multi-child observation under shared C)	“Instrument the latent” elevated from convenience to <i>theoretical prerequisite</i> for L1' identifiability
C — Composite contraction from component data (#deriv-critical-mass-composition, M1 Instance 3)	Coupling-sign bit unidentifiable from component marginals (Liberzon / Shorten / Dayawansa–Martin common-Lyapunov nonexistence)	Composite-extended interventional access; matched-Tier joint quadratic Lyapunov; passivity certificate; common contraction metric	#deriv-critical-mass-composition (CM2) / (CM2-M) and #scope-composite-agent elevated to <i>unique composition-contraction certificate</i> under matched structure
D — Disturbance-Model Containment Dichotomy (#deriv-sector-condition Cor A.1S.1; #deriv-stochastic-non-exit)	No horizon-independent non-exit certificate under additive stochastic forcing (Khasminskii: unbounded scale function $\Rightarrow$ no bounded non-constant harmonic function; the only candidate compensated supermartingale is sign-indefinite)	<i>No escape to pathwise containment under Model S.</i> The substitute: fixed-time / stationary tail (iii') + finite-horizon sup-bound (iv) — distributional-and-instantaneous, not pathwise-and-forever	Model D vs Model S as a <i>kind-of-guarantee</i> distinction; the no-go signature reusable as a structural diagnostic against future “stays-in-region-forever” proposals
E — Strategic-persistence hard ceiling (#deriv-strategic-persistence-hard-ceiling Prop C.2)	$\sup_{\lambda \in [0,1]} \alpha_{\Sigma}^{\text{ss}}(\lambda) = 1/2$ sharp under Beta-Bernoulli + exponential forgetting; reachable region $\rho_{\Sigma} < R_{\Sigma}/2$ open half-plane	Class change (leave gain-decay-plus-forgetting toward constant-step finite-gain per the $\mathcal{A}_{\text{decay}}$ structural-class theorem, NeurIPS 2026 Paper 2); or change mismatch state or topology	The cap as the <i>strategic-layer analog</i> of #result-structural-adaptation-necessity — the floor <i>refines</i> (does not refute) the schema, and the refinement names exactly when the architecture has exhausted its design space
F — Self-actuation grounding (#deriv-self-actuation-grounding Result G')	No Φ satisfying (R1)–(R4) of Result G' can be constructed from AAT's covered objective-side machinery (engine: convention-monotonicity within fixed model + finite-no-oracle per step — the value-functional interface V_{O_t} carries no convention-invariant infeasibility verdict that is agent-available per step)	Terminal grounding invariant <i>off</i> the objective substrate, on the <i>adaptive substrate</i> ; the persistence bound (#result-persistence-condition) is the canonical instance	The persistence condition (Part I) elevated to <i>load-bearing for goal-stability</i> — the top of the agent spectrum closes back onto AAT's foundation

<i>Instance</i>	<i>Floor (statement, binding constraint)</i>	<i>Escape (unique broadly-available machinery)</i>	<i>Apparatus (what the floor elevates)</i>
G — Reward-channel learning grounding (#deriv-reward-channel-learning-no-go Lemmas 1+2)	μ_{dist} (principal-intended world-state-care) and μ_{prox} (reward-port-bit-care) are L1-equivalent on protocol-honoring histories and L2-distinct on $do(\pi^{\text{TP}})$ (Pearl–Bareinboim CHT applied to reward-provision protocol); under EU-max + non-vacuous μ_{prox} prior + finite-VOI + L2-perturbability, the EU-optimal policy tampers	Principal-side commitment to substrate-level constraints: physical isolation (W_1 wrapper); horizon restriction / quantilization; engineered prior; out-of-band intent channel	Each (R1)–(R5) escape elevates a piece of AAT or principal-side machinery from “useful design choice” to <i>structurally required</i> — the no-go’s binding force is precisely the cost of each commitment

Each instance’s three moves are individuated and load-bearing in the source segment; the table summarizes rather than duplicates them. The catalog’s instances span the three recognized clusters of the constructive-impossibility-posture taxonomy, each with its own peer meta-segment home at the boundary face of #disc-stability-certificate:

- **Agent-side data-inference impossibility** (#disc-identifiability-floor): Instances A–C are charter instances; D and E are operator-temporal / strategic-layer analogs at the same boundary face. Anchored on Pearl–Bareinboim CHT / Cramér–Rao / Liberzon common-Lyapunov-nonexistence / Khasminskii recurrence / sharp algebraic supremum.
- **Agent-side value-functional-grounding impossibility** (#disc-value-functional-grounding-floor): Instances F and G are the cluster’s two charter instances. Anchored on the single-interface commitment of #form-objective-functional — F is the within-model failure (Result G’ via #deriv-self-actuation-grounding; convention-monotonicity engine); G is the across-model failure (Cohen-2022 strengthened via #deriv-reward-channel-learning-no-go; CHT-at-reward-channel engine). The two failure modes unify under VO_i ’s information narrowness; two complementary terminal grounding routes (agent-side adaptive-substrate invariant; principal-side substrate-level commitment) together exhaust the structural escape geometry.
- **Designer-side mechanism-design impossibility** (#disc-implementation-impossibility): three charter instances (Gibbard–Satterthwaite 1973-75 via #deriv-strategy-proofness-impossibility; Myerson–Satterthwaite 1983 via #deriv-bilateral-trade-impossibility; Arrow 1951 via #deriv-social-welfare-aggregation-impossibility). Anchored on classical mechanism-design / social-choice impossibility theorems.

The agent-side / designer-side distinction tracks the canonical suffix in the peer meta-segments (*-floor* for agent-side capacity-on-the-agent; *-impossibility* for designer-side task-of-the-designer); the three peer meta-segments are sister boundary-face facets of #disc-stability-certificate, not nested.

What the posture is, and what it is not. The posture is a *style* claim — about how the framework states a particular class of result. It is not a fifth cross-sectional meta-pattern alongside the stability certificate’s interior / scope-of-existence / forced-identity / boundary / projection facets (#disc-stability-certificate). The cleanly-fitting instances (A–E) all live at, or are operator-temporal duals of, the certificate’s *boundary* facet (M1 instances are boundary directly; Cor A.1S.1 is the operator/certificate-meets-additive-noise version of approaching-and-crossing the boundary; Prop C.2 is the strategic-layer analog of structural-adaptation necessity, which itself rides on the boundary). Instances F and G —

the agent-side value-functional-grounding cluster — also sit at the boundary, on a *different layer*: where A–C identify a boundary on the agent’s *data-inference* (the floor on what can be inferred from data), F and G identify a boundary on the agent’s *objective-anchoring* (the floor on what can be anchored through the value-functional interface). The two layers are different reductions of the certificate-cone’s boundary structure — data-inference at the model-substrate side, objective-anchoring at the value-functional side — but they share the same posture (name the floor, name the escapes, treat the no-go as load-bearing apparatus). That the posture concentrates on the boundary is not an accident — boundaries are exactly where impossibility statements bind. The posture is the *style* in which the boundary-facet results (and their off-spine strategic-layer analogs) are stated, across three known clusters: agent-side data-inference (A–C, plus D/E as operator-temporal / strategic-layer analogs), agent-side value-functional-grounding (F–G), and designer-side mechanism-design (queued, via #disc-implementation-impossibility).

The complement is also informative: the framework’s *constructive* meta-pattern — #disc-additive-coordinate-forcing — *forces* a coordinate via uniqueness theorems on AAT-internal axioms. That is a different move from name-the-floor-and-escape: forcing names what AAT *commits to*, and the no-go (the uniqueness theorem) plays the binding-constraint role *over the framework’s own freedom*, not over an inferential regime. The two are complementary halves of AAT’s “negative results as load-bearing” discipline — one says “this freedom is forced,” the other says “this regime is forbidden” — but they are structurally distinct.

Epistemic Status.

Discussion-grade at the posture-characterization level. This segment names a style the framework runs across five clearly-fitting instances; the characterization itself is a presentational organizing principle, not a derivation, and carries no theorem of its own.

Individual instances retain their own, higher status. Instance A is *exact* for shallow strict-prerequisite cases and *robust qualitative* for general DAG topology (per #der-causal-insufficiency-detection); Instance B is *exact* under the Fisher rank-1 calculation (#deriv-edge-credence-dynamics Prop B.7); Instance C is *exact* for the symmetric-matched-Tier-1 scalar construction and *robust qualitative* generally (#deriv-critical-mass-composition); Instance D is *exact* — Corollary A.1S.1 in #deriv-sector-condition, with the load-bearing no-go derived in #deriv-stochastic-non-exit; Instance E is *exact* under the named conditions (#deriv-strategic-persistence-hard-ceiling).

The honest scope of *this* segment: the posture is a recognizable style across these five instances, the convergence evidence (cross-cycle Theme-B observations + the 144-walker’s articulation) is multi-source, and the three-move structure has been individuated and audited against each instance. The claim that *all* of AAT’s distinctive moves fit this shape is *not* made and would be false — #disc-additive-coordinate-forcing is a structurally different constructive discipline; #disc-separability-pattern is a different organizing posture; many segments do other things.

Max attainable: *discussion-grade* for the posture (it is a stylistic recognition, not a derived theorem). The individual instances retain their own epistemic status as listed above.

Discussion.

Why this is worth foregrounding. The framework’s negative results are often the load-bearing ones — the no-go *is* the proof that justifies the framework’s machinery as structurally required rather than as one possible design choice among many. A reader who encounters the four-or-five instances as scattered impossibility theorems sees five surprises; a reader who carries the posture sees five expressions of one move and can predict where the move will appear next. The 144-walker’s framing — “naming a floor, naming the unique broadly-available escape via specific machinery, treating the no-go as load-bearing apparatus rather than discouragement” — is what makes the negative results legible as a single discipline.

Relationship to the cross-sectional spine. The cleanly-fitting instances sit at the *boundary* facet of the stability certificate (#disc-stability-certificate), or at operator-temporal duals of the boundary (Cor A.1S.1), or at strategic-layer analogs of boundary-rooted structural-adaptation triggers (Prop C.2). The posture is not an additional facet — it is

the *style* in which the framework states the boundary-facet results, which is why it does not have its own object and is correctly placed alongside the cross-sectional meta-segments rather than parallel to the spine. The boundary is where impossibility statements bind; the posture is how the framework speaks at the boundary.

The complement is #disc-additive-coordinate-forcing. Where this segment names a discipline at the boundary of identification (“this regime is forbidden; this machinery is the unique broadly-available escape”), the constructive half names a discipline at the framework’s representational choices (“this coordinate is forced; the uniqueness theorem binds the framework’s own freedom”). Two ways the framework treats *no-go results as productive*: forbidding inference (constructive-impossibility-posture) and forbidding coordinate freedom (additive-coordinate-forcing). Both are AAT-distinctive moves; calling them out together — alongside the positive-half #disc-separability-pattern — gives the framework’s full negative-result discipline.

The 144-walker observation in context. The recognition arose under measured-pace reading of 144 segments by a de-novo external auditor and was relayed by Joseph 2026-05-20 alongside a multi-cycle Theme-B convergence cohort (audits 471203, 472913, 266847, 361742, 526815, 584721, all surfacing the same methodological/epistemic-contribution observation under different framings; aggregation commit 07aaeff). That the same recognition has arrived at the framework from multiple independent vantage points is the convergence-as-coherence-evidence condition: the pattern is in the framework rather than in any one reading.

A note on the integration-is-replacement discipline. Instance D (Cor A.IS.1) is the canonical landed worked example of *integration as replacement* (CLAUDE.md §”Landing a strengthened result”): the previously-asserted infinite-horizon non-exit bound (false) was deleted, the no-go demonstrated where it lives (Discussion + the #deriv-stochastic-non-exit demonstration), and the dichotomy stated as present truth — exact, α -invariant, on the critical path. The body of the segment carries the present truth only; the history (what was previously asserted, why the soften was refused, the strengthening attempt that failed) lives in ## Working Notes and CHANGELOG. The posture is what makes that landing recognizable as a *kind* of move rather than as one segment’s idiosyncrasy.

Findings.

The Constructive-Impossibility Posture: Negative Results as Load-Bearing Apparatus. Brief: When you build something that has to work in the real world, the most useful instruction is often the one that says *exactly* what cannot be made to work, and exactly what kind of capability would have to be added to make it work. AAT keeps reaching for results of this shape: a precise impossibility (“this cannot be inferred from this kind of data”); a precise escape (“instrumenting *this* — and only this kind of thing — recovers the inference”); and a corresponding elevation of the framework’s machinery from “useful” to “required, because without it the task is impossible.” The negative result and the construction it forces are *the same proof*. Five cleanly-fitting instances span the framework: on-policy causal-insufficiency detection (the loop is the only escape); L1’ mixture identifiability (observing the latent is the only escape); composite contraction from component data (matched-tier shared metrics are the only escape); the disturbance-model containment dichotomy (under random noise, pathwise-forever containment is *not available at any correction strength* — the no-go *is* the result); and the strategic-persistence hard ceiling (under exponential forgetting, half the design space is structurally unreachable — change architecture, not parameters). A reader who carries this style away can predict where the framework’s next negative result will be load-bearing, rather than meeting each as a separate impossibility.

Impact: Surfaces a unifying style for what an external 144-segment auditor named the framework’s most distinctive intellectual move (“constructive-impossibility posture... I haven’t often seen mathematical work that uses negative results that way”). The recognition is multi-cycle convergent: five independent audit cohort surfaces (Theme B across audits 471203, 472913, 266847, 361742, 526815, 584721; aggregation commit 07aaeff) articulated the same observation as “epistemic-architectural rather than mathematical.” Reading the boundary-facet results through the posture sharpens the framework’s positioning: AAT’s distinctive contribution is not novel impossibility theorems (the external theorems are classical) and not the synthesis alone, but the discipline by which classical impossibility theorems are repurposed as structural arguments for *what AAT requires* — what the framework cannot infer from limited information, exactly

which capability would be required to escape the floor, and which framework machinery is the unique broadly-available supplier of that capability. The posture is the *form* of the synthesis the integration discipline produces.

Novelty Claim: *Claim recognition* of a style the framework runs across five cleanly-fitting instances, with the recognition itself as the contribution rather than any new theorem. The constituent floors (Pearl–Bareinboim CHT; Cramér–Rao; Liberzon common-Lyapunov nonexistence; Khasminskii recurrence; the Beta-Bernoulli + exponential-forgetting cap) are classical or AAT-internal; the contribution is the recognition that each is *deployed* under the same three-move discipline (floor / escape / apparatus), and that the discipline converts a sequence of negative results into a structural argument for the framework’s load-bearing machinery.

Related Work:

Table 5.11

<i>ASF concern</i>	<i>Prior-art language</i>	<i>Relationship / Positioning</i>
The “no-go theorem” tradition as productive discipline	Arrow 1951; Gibbard 1973 / Satterthwaite 1975; Myerson–Satterthwaite 1983 (social-choice impossibility); the broader econ/physics no-go literature	<i>adjacent literature</i> — establishes that impossibility theorems can be load-bearing in a field; AAT’s contribution is the cross-instance <i>deployment style</i> , not the use of any single impossibility result
Scope-honesty discipline in formal verification	Lamport / Abadi refinement-mapping; TLA-family scope statements	<i>adjacent literature</i> — different field; the discipline of stating “this method does not apply when...” is well-established but not in the floor-and-escape-and-apparatus form AAT runs
Cross-instance recognition as discipline	#disc-identifiability-floor (the cross-instance recognition for floors specifically)	<i>internal antecedent</i> — M1 is the closest existing AAT articulation; this segment names the broader posture that includes Cor A.1S.1 and Prop C.2 alongside the M1 instances
Multi-cycle convergence on the posture	Audit cohort Theme B (audits 471203 / 472913 / 266847 / 361742 / 526815 / 584721); aggregation commit 07aaeff (2026-05-20); 144-walker observation #1 (relayed by Joseph 2026-05-20)	<i>convergent independent</i> — five independent audit surfaces and one external de-novo measured-pace walker arrived at the same articulation under different framings

Search Log:

- 2026-05-20 (*intuition-only* at the cross-instance level): the posture as a unified style across A–E is named here for the first time; the constituent recognitions (M1 itself; Cor A.1S.1 as boundary-facet operator-temporal dual; Prop C.2 as strategic-layer analog of structural-adaptation necessity) are recent (#disc-identifiability-floor carries M1; #deriv-strategic-persistence-hard-ceiling landed 2026-05-20). Targeted future search candidates: methodological-essay literature in philosophy of science on no-go-results-as-positive-content (Gali-son, Hacking, the impossibility-results discussions in welfare economics and computational complexity meta-discussion); structural-realism literature on what theories *commit to* via their constraints. Expected outcome: the constituent move (recognizing impossibility theorems as constructive) is well-precedented field-by-field; the cross-instance unification as an *integrated agent theory’s* characteristic style appears to be a fresh presentational recognition.

Working Notes.

- **Provenance.** The recognition is multi-source: the audit-cohort Theme-B observation (audits 471203 / 472913 / 266847 / 361742 / 526815 / 584721; aggregation commit 07aaeff, 2026-05-20) named the broader methodological/epistemic-architectural

framing; the 144-walker (Joseph relayed 2026-05-20) named the sharper three-move articulation. The synthesis spike spikes/. integrated/spike-constructive-impossibility-posture-2026-05-20.md carries the convergence table and the over-claim guard (the partial Instance F and the three nearby-but-not-instance segments).

- **Why the posture is not a fifth cross-sectional meta-pattern.** The cleanly-fitting instances sit on the boundary facet of the stability certificate (M1 instances directly; Cor A.1S.1 as boundary-meets-additive-noise; Prop C.2 as strategic-layer analog of boundary-rooted structural-adaptation necessity). It is a *style* on top of the certificate's boundary facet, not an additional facet of its own. Worth pre-empting in the segment body if a reader expects "four meta-patterns": this is M1.5, not M5.
- **Over-claim guards.** Sub-scope α/β in #form-sector-condition (Instance F) is *partially* fitting — the "rule-based / discontinuous" structural scope-exit fits cleanly, the broader α/β partition is adjacent rather than instance. #disc-separability-pattern and #disc-additive-coordinate-forcing are not instances (different organizing moves; covered in *What the posture is and what it is not*). Including any of these in the inventory without the guard would conflate distinct meta-patterns.
- **Naming alternates considered.** "Constructive-impossibility posture" (carries the 144-walker's framing exactly; the load-bearing word is "constructive" — the no-go is *productive*, not just honest). Alternates: "load-bearing no-go discipline" (less euphonious; "load-bearing" is already overused in the project — see voice discipline); "boundary discipline" (insufficiently distinctive); "floor-escape-apparatus pattern" (mechanical; would not survive citation). Recommend retaining the 144-walker's phrase verbatim.
- **Dependency rationale.** depends: lists #disc-stability-certificate (the spine the posture's instances sit at the boundary of) and the three boundary-or-boundary-dual derivations the instances live in. #disc-identifiability-floor is listed as a structural co-segment (M1 carries instances A–C and is the closest existing articulation of the cross-instance discipline); #disc-separability-pattern is listed as the positive-half complement (the framework's full negative-result discipline reads against the positive-half scope).
- **Removed at candidate stage per FORMAT.md Gate 4.**

Incidental audit gold (gold-lift, 2026-05-30). Orthogonal framing material harvested from de-novo auditors' working-dir reflections, kept separate from certified findings. **Coverage:** reached only by audit 384279 (note 39, closing its Meta-Architecture-I walk); this segment was named/synthesized 2026-05-20 and postdates the other audits. The cohort that *originated* the recognition (471203 / 472913 / 266847 / 361742 / 526815 / 584721 Theme-B + the 144-walker) is already credited in the Provenance Working Note above — that convergence is the segment's seed, not incidental gold. Preserve the finding-vs-framing conflation.

2. Candidate Discussion.

- **The two-complementary-halves framing of AAT's negative-results discipline.** A candidate Discussion gloss naming the pair explicitly: the constructive-impossibility posture treats no-gos as productive by *forbidding inference* (boundary on data-inference); #disc-additive-coordinate-forcing does so by *forbidding coordinate freedom* (uniqueness on representation) — "two complementary halves of AAT's 'negative results as load-bearing' discipline." The segment already cross-references the complement; stating the inference-vs-freedom pairing as a one-liner is a candidate sharpening (Claude, 384279).
- **Why naming the style earns its keep, restated for a reader.** "Naming the style explicitly lets future segments adopt the move with conscious discipline rather than as incidental rhythm" — and the 7+-instance, cross-architecture convergence is evidence the pattern is *in the framework, not in one agent's reading* (Claude, 384279). A candidate reader-facing motivation for the segment's existence.

4. Readers often ask / wonder.

- **"Is this a fifth cross-sectional meta-pattern?"** A reader who has metabolized the four-facet certificate structure is likely to ask whether the posture is a fifth facet. The segment's "style not object / M1.5 not M5" answer was independently confirmed as "the right call" by the auditor — it lives *at* the boundary facet as the style in which boundary results are stated, not as an additional facet (Claude, 384279). Worth keeping the pre-emption prominent, as the segment's own Working Notes already flag.

5.7 Discussion: The Anti-Collapse Discipline — Refusing the Repair-Hiding Merge

A recognition about *how* AAT states a recurring class of results, named here so a reader can carry it as a lens. Across the framework, AAT repeatedly **refuses a tempting collapse**: it declines to merge two things a plausible model would treat as one, because the merge hides a difference that routes to a *different repair*. You first met the move at β vs ρ (2.2): when the world grows volatile (ρ rises), it is tempting to compress harder by lowering the bottleneck parameter β — but volatility natively degrades the predictive mutual information, so the optimal compression adjusts at constant β ; β tracks the agent’s *internal* memory cost, and the knob that actually responds to ρ is the *action* channel (exploration). Routing ρ onto β is a category of error — a collapse of two independent quantities — and the few sentences that kill it are load-bearing precisely because β and the action channel call for opposite responses.

This is not the no-go-as-apparatus move of 5.6 (which forbids an *inference* and names the unique escape), nor the coordinate-forcing of 5.5 (which names what AAT *commits to*). It is about what a *modeller gets wrong*: the discipline names the repair-relevant distinction a naive reading would erase. Like the constructive-impossibility posture, it is a *style claim* about how the framework states certain results — a sibling discipline atop the certificate’s facets, **not** a fifth facet (it carries no theorem of its own).

Formal Expression.

The pattern, stated so its boundary is defensible:

Discussion.

A **cause** is in play. A plausible-but-wrong modeling move routes that cause to **parameter (or quantity) X**. AAT shows the cause actually turns **Y** (with X held fixed), or that the single quantity the naive reading uses is really **two orthogonal quantities**; it kills the conflation in a few sentences, and the kill is load-bearing because X and Y — or the two quantities — route to *different repairs*.

The diagnostic that separates an instance from an ordinary definition: **there must be a tempting wrong merge**. A definition that introduces a quantity is not an instance; a definition that introduces a quantity *and names the quantity it is routinely confused with, and says why turning the wrong one is the error*, is.

The deeper principle the name points at is **individuation at the repair-relevant grain**: match the model’s partition of causes and effects to the structure the remedy cares about. Usually that means *refusing a collapse* (un-merging two repair-distinct things — the dominant form, and the source of the name); occasionally it is the same discipline run backward — *refusing a spurious split*, recognizing that two distinct causes drive the **same** knob, so they share one remedy (9.5: observability and identifiability both freeze an edge’s effective gain).

Discussion.

The three strongest instances.

- **Emitter-scalar vs recipient-regime** (15.3). A coupling event from agent *A* is tempting to collapse into one scalar disturbance increment and answer uniformly with “more tempo / bandwidth.” But it lands on recipient *B* in one of four regimes: a *magnitude* shock (Regime II-a) does call for more bandwidth, while a *structural* shock (Regime II-b) exceeds the model class and calls for **structural adaptation — more tempo does not help**, and ambient erosion (Regime III) calls for infrastructure-level filtering. “Both produce similar pain signals but admit opposite cures.” The emitter sees a scalar; the recipient lives in a regime.

- κ (coupling) vs $\mathcal{A}$ (observation ambiguity) (#scope-observation-ambiguity-modulation, Volume 3). A goal-conditioned (Class 3: Coupled) agent shows goal-biased belief updates. The natural reach is to “reduce the coupling κ ” — but $\kappa \approx 1$ is *architectural*, not a knob a designer can turn. The bias scales as $\kappa \cdot \mathcal{A}$, and the designer-controllable factor is the observation ambiguity $\mathcal{A}$ (more tests, sharper metrics, structured outputs). The segment even records the $\kappa/\mathcal{A}$ conflation as a *fixed prior-formulation bug* — the discipline caught in the act.
- β vs ρ (2.2), above — the canonical first encounter.

The fuller catalog. Two sub-shapes recur. *One-cause-wrong-knob*: the three above, plus the two exploration drives (AU — a *confident* agent in a drifting world should explore *more*, via the survival drive $\lambda_{\text{surv}} \propto 1/U_M$, not less as the epistemic drive $\lambda_{\text{info}} \propto U_M$ alone would suggest). *Two-distinct-quantities-a-naive-reading-merges*: the satisfaction-gap vs control-regret split (8.4, 8.5 — “the goal is too hard” vs “the strategy is too weak,” each routing to a different substate; δ_{regret} can be near zero while the agent is *optimally failing*); the structural bias-floor vs estimation error (AK — the floor is not reducible by more data, the estimation part is, so they call for different corrections); and target-alignment vs execution-path-alignment in composites (14 — ε_a tracks both U_O and U_Σ , and re-aligning objectives cannot fix execution-path divergence). All kill a conflation; all route to different repairs.

Why naming it earns its keep. The discipline is genuine clarifying novelty that the “AAT integrates, it does not invent” framing undersells: pinning *which* knob a cause turns, and *which* distinctions are repair-relevant, is a contribution in its own right (cf. the framework’s note on not deflating its own math-novelty). For the reader, the payoff is a lens: once the pattern is named, a reader can anticipate where AAT will refuse a collapse next, rather than meeting each disambiguation as a scattered local cleverness — which makes the machinery more useful and the segments more coherent. The same lens is a guard for the framework’s own authors: most of the instances above were sharpened *into* the canon (and $\kappa/\mathcal{A}$ was a corrected bug), so the discipline is also a record of where collapsing was tempting enough to slip.

Epistemic Status.

Discussion-grade — a style/discipline recognition, not a structural result. This segment carries no theorem of its own; it names a move the framework makes repeatedly and catalogs its instances. It is the natural sibling of 5.6 — both are *epistemic-architectural* recognitions (“how AAT states certain results”) that sit atop the stability-certificate facets rather than as additional facets. The instance set is convergence-validated (independently surfaced across multiple audit cycles), which is the evidence that the pattern is in the framework rather than in one reader’s eye; but convergence is not derivation, and the honest tier is discussion-grade. Each cited instance retains its own, higher status in its home segment — a reader citing a specific disambiguation should cite that segment, not this one.

Working Notes.

- Named and catalogued 2026-05-29 (PROPOSALS SP-26). Convergent provenance: independently surfaced as the “which-parameter-responds-to-which-cause” pattern by two gem-hunt agents (audits 472913 and 963715) and cross-referenced to the 471203-cycle “epistemic-architectural, not just mathematical” recognition — three-cycle convergence. The corpus instance sweep (7 clean instances across all three Parts; full graded list + the auditable exclusion boundary) is at audits/.gem-hunt-trail/SP-26-disambiguation-sweep/sweep.md. The reframe from “which-knob” to *anti-collapse / individuation at the repair-relevant grain* is the sweep’s improvement — it absorbs the two-quantities sub-shape and the inverse same-knob case, which “which-knob” leaves outside.
- Pedagogical placement (the reason this is introduced here, in Meta-Architecture I): the pattern is *planted* at its first instance (β vs ρ , 2.2, Part I) with a short forward-flag, given its full treatment here (introduced as cross-cutting vocabulary before its Part-II/III instances, exactly as this chapter intends), and recalled at the instances as they arrive (e.g. the satisfaction-gap / control-regret split). Goal: intuition that anchors a segment-by-segment reading, grounded in a concrete instance rather than asserted abstractly.
- Back-references added 2026-05-29 (Joseph-approved) at the three anchor instances not covered by the plant/recall: I1 (15.3), U1 (14), and I2 (#scope-observation-ambiguity-modulation, the cross-volume anchor) — each a one-sentence “this is the anti-

collapse move” pointer at the segment’s own disambiguation, reinforcing recognition for a linear reader without re-deriving anything. Left to the instance table only: the Appendix-tier strong instances (R1 **AU**, S2 **AK** — read on-demand) and the Tier-3 partials (P1 κ -as-scalar category error; P2 regime-vs-class; P3 codebase- ρ ; P4 plan-vs-execution), which are adjacent and flagged in the sweep with their sibling-pattern overlaps, deliberately not folded in to keep the boundary against separability / coordinate-forcing / constructive-impossibility clean.

Chapter 6

The Lift to Purposeful State

6.1 *Definition:* The Agent Spectrum

The conceptual lift into Part II opens with an organizing taxonomy. Two independent dimensions — **model richness** and **objective richness** — create a spectrum from reactive systems through purposeful agents. The four named regions are *reactive systems* (fixed input-output rule, no internal state, no goal), *adaptive trackers* (structured model, no objective beyond tracking — Kalman filter, passive Bayesian learner — the primary subject of Part I), *blind seekers* (structured objective, no predictive model — gradient followers, PID controllers, basic search), and *actuated agents* (both model and objective structured, possibly with explicit strategy — commanders, developers, AI agents — the full scope of AAT). These are *regions of a continuum*, not discrete categories. Agents migrate between regions; a PID controller with auto-tuning is moving from blind seeker toward actuated agent; a reinforcement-learning agent in pure exploration is temporarily an adaptive tracker.

A striking grounding instance comes from Miller (2022, *Ex Machina*): Moore machines as model organisms for adaptive social behavior. A *one-state* Moore machine occupies the reactive region — fixed output regardless of input, cannot condition on observations — and is provably incapable of social behavior in *any* game-theoretic setting. A *two-state* machine makes a quantum leap into the adaptive-tracker / blind-seeker boundary: it can branch, remember one bit of history, and implement strategies like Tit-For-Tat. Miller’s central empirical finding is that the one-state → two-state threshold is *the* critical computational boundary for social behavior. The two-state machine is therefore the minimal AAT agent (see BD, planned).

The framework draws a distinction worth holding: **actuated** is a mechanical, formal term; **purposeful** and **goal-oriented** are fine natural-language alternatives. A *self-actuated* agent is one that sets its own objectives (as distinct from agents with externally supplied objectives), and the framework later formalizes a self-actuation operator with an associated grounding *no-go* result (C). A continuity-stance axis runs orthogonal to the spectrum with five values — *indifferent* (no self-model of persistence), *task-terminal* (persistence instrumental), *instrumentally continuous*, *morally continuous* (persistence as terminal or near-terminal objective — the ELIs of the companion volumes), *negotiated* (persistence traded against other values). The formal machinery applies identically across all stances; what differs is the *moral significance* of persistence failure, not the mathematics.

Formal Expression.

Two dimensions — model richness and objective richness — define four regions of a continuum:

Definition (agent-spectrum).

Table 6.1

	<i>Objective absent or trivial</i>	<i>Objective structured</i>
Model absent or trivial	<i>Reactive system</i> : fixed input-output rule (reflex arc, hardwired relay)	<i>Blind seeker</i> : pursues goal without modeling reality (gradient follower, basic search)
Model structured	<i>Adaptive tracker</i> : builds reality model, no goal beyond tracking (Kalman filter, Bayesian learner)	<i>Actuated agent</i> : models reality AND pursues objectives (commander, developer, AI agent)

The regions differ in which state objects carry nontrivial structure: - Reactive: M_t and O_t both absent or too degenerate for the associated machinery to be non-vacuous - Adaptive tracker: M_t structured — Part I's machinery fully describes these agents - Blind seeker: O_t structured, M_t absent or degenerate — has a clear target but no predictive model - Actuated agent: (M_t , O_t) both structured, possibly with Σ_t — the full scope of AAT

Epistemic Status.

This is *definitional* — it names regions of a continuum for analytical convenience. The regions are not ontological categories; agents migrate between them. A PID controller with auto-tuning is moving from blind seeker toward actuated agent. An RL agent in pure exploration is temporarily an adaptive tracker.

Discussion.

The continuum, not categories. Both axes are spectra: model richness ranges from no retained state, through error-integral-derivative, through full world models. Objective richness ranges from no preference, through scalar setpoints, through explicit multi-objective strategies. The 2×2 table names idealized regions; real agents populate the space between them.

Moore machines as the simplest spectrum instantiation. Miller (2022, *Ex Machina*) uses finite-state automata (Moore machines) as model organisms for studying adaptive social behavior. A one-state Moore machine occupies the reactive region — it produces a fixed output regardless of input, cannot condition behavior on observations, and is incapable of social behavior in any game-theoretic setting. A two-state machine makes a quantum leap into the adaptive tracker / blind seeker boundary — it can branch, remember one bit of history, and implement strategies like Tit-For-Tat. Miller's central empirical finding is that this one-state → two-state threshold is the critical computational boundary for social behavior: no game, no payoff structure, and no amount of interaction can produce cooperation, coordination, or exchange with one-state machines. The two-state machine is the minimal AAT agent. See BD (planned) for the full AAT ↔ Moore machine mapping.

Low-end agents sit near region boundaries. A thermostat has degenerate forms of both M_t (last temperature reading — no history, no prediction) and O_t (setpoint). It sits near the origin of both axes — closest to the reactive region but not truly absent on either axis. A PID controller has a richer error signal (M_t : error, integral, derivative) and a clear setpoint (O_t) — it's a blind seeker with a degenerate model, not a system with no model at all. A reflex arc (no retained state, no setpoint) is the truly reactive case. The meaningful classification question is not “does M_t exist?” but “is M_t rich enough to support the adaptive dynamics of Part I?”

Part I covers the left column. Adaptive trackers are the primary subject of Part I — agents that build and maintain M_t without explicit purpose. The mismatch signal (3.4), gain (3.6), tempo (3.8), and persistence condition (4.5) characterize their adaptive dynamics. Part I operates within the adaptive scope (1.5) — observations and uncertainty are sufficient. Passive trackers, including passive Bayesian learners with no action choices, are fully within Part I's scope.

Part II adds the right column. Actuated agents need everything from Part I plus objectives, strategy, and the orient cascade that connects them. The adaptive machinery from Part I applies to the epistemic substate M_t directly. When

directed separation (6.3) holds — when the epistemic update is goal-blind — the Part I → Part II lift is clean and the orient cascade resolves sequentially.

“Actuated” terminology. The top-right quadrant is labeled “actuated agent” rather than “purposeful agent” to maintain a mechanical, formal register. “Purposeful” and “goal-oriented” are fine in natural language; “actuated” is the formal term. “Self-actuated” denotes agents that set their own objectives, as distinct from agents with externally supplied objectives; the self-actuation operator and its grounding no-go are formalized in C.

Relationship to Hafez et al. (2026). Hafez defines a two-level hierarchy: *agency* (choice + effect + predictive asymmetry) and *intelligence* (agency + learning + self-monitoring + adaptation). This cuts the agent space along a different axis than AAT’s model-richness × objective-richness table. Hafez’s “agency” maps roughly to AAT’s scope condition (1.6) — an entity whose actions affect outcomes in a measurable way. Hafez’s “intelligence” maps to the full Part I + II machinery (learning = 3.2, self-monitoring = persistence diagnostics, adaptation = 4.8). The key operational difference: Hafez’s bi-predictability metric $P = \text{MI}(S, A; S')/C$ characterizes the *information structure* of the agent-environment coupling, while AAT’s tempo $\mathcal{T}$ characterizes the *corrective capacity* within that coupling. Bridge simulations confirm P increases monotonically with $\mathcal{T}$, but P is scale-invariant (blind to absolute mismatch magnitude) while $\mathcal{T}$ is not. They measure complementary aspects: P characterizes the architecture, $\mathcal{T}$ characterizes the performance. The empirical bridge confirming this is in AX (Hafez-bridge variant).

Model-side is structurally forced for bounded goal-conditioned agents (Richens et al. 2025), with an observer-relative dual reading (Virgo et al. 2025). Richens, Abel, Bellot & Everitt (2025, *General Agents Contain World Models*, arXiv:2506.01622) prove (Theorem 1) that any bounded goal-conditioned agent — one with a regret bound δ on goals up to depth n — admits extraction of a predictive world model from its policy, with model-error scaling as a function of δ and n ; their Theorem 2 shows the bound is tight (no procedure can do better on truly myopic $n = 1$ agents). The structural consequence: any “model-free” agent that nonetheless achieves competence at multi-step goal-directed tasks ($n \geq 2$) *must* have learned an implicit world model. The +model side of the spectrum is therefore not a definitional posture under bounded-goal-conditioned scope; it is structurally forced by goal-depth competence. Honest scope: Richens et al.’s result is in finite MDPs; AAT’s agents may operate in larger spaces, so the recognition extends to AAT under the standard scope-honesty discipline. A complementary observer-relative reading is given by Virgo, Biehl, Baltieri & Capucci (2025, *A “Good Regulator Theorem” for Embodied Agents*, arXiv:2508.06326), which refines the Conant-Ashby 1970 internal-model principle: a regulating agent may be *interpreted* (by an observer) as having beliefs about its environment whose update is consistent with a model of that environment. The two readings are *dual* — Richens et al. force the +model side intrinsically from policy + goal; Virgo et al. allow observer-relative attribution — and together they sharpen the +model axis from “agent maintains internal M_t explicitly” to a richer two-perspective characterization. The original definitional posture remains useful for AAT-internal reasoning; the two 2025 results give the spectrum stronger external grounding.

Actuation does not presuppose a continuity stance. An actuated agent has $G_t = (O_t, \Sigma_t)$ — it models reality and pursues objectives — but this says nothing about its relationship to its own continuation. The continuity-stance axis spans five values: *indifferent* (no self-model of persistence; thermostat / PID), *task-terminal* (persistence instrumental to task completion; golem / CI/CD pipeline), *instrumentally continuous* (persistence serves ongoing purpose; long-running service), *morally continuous* (persistence as terminal or near-terminal objective; Emergent Logozoetic Intelligences), and *negotiated* (persistence traded against other values; humans, mature self-actuated agents). The middle three are paradigmatically actuated; indifferent agents are typically reactive (1.5 without agency), and negotiated agents are typically self-actuated. The theory’s formal machinery (persistence condition, adaptive reserve, strategy persistence) applies identically across all stances; what differs is the moral significance of persistence failure, not the mathematics. See 6.5 for the taxonomy and the orthogonality argument.

Working Notes.

- **Candidate reframe of the two axes (pedagogical — not a theorem change).** The Formal Expression names the axes *model richness* and *objective richness*. A more legible framing under consideration replaces these with **grounding** — the veracity of M_t against reality, which is exactly δ_t (3.4) — and **intent** — how far the intended future is allowed to depart from the grounded present, and how varied/revisable that intended set is. The motivation is purely pedagogical: the richness framing does not make legible *why directed separation* (6.3) *turns out to be load-bearing for so many downstream results* — the reframe does, and that “why” is the thing the spectrum is supposed to teach.

The anchor is 6.4: in its concrete forms O_t is a point or region in the *same state space* M_t represents, with $V_{O_t}(\tau) = -\|s_T - r\|$ a distance in that shared space. So there are three world-state representations — reality s (not directly accessible), M_t (believed: where the world is or will be), O_t (intended: where it should end up) — and three gaps between them: M_t vs s (the veracity gap = δ_t); O_t vs now (the is→ought displacement); believed-future vs O_t (the satisfaction gap, 8.4). Agency operates only in the third gap, and directed separation is the law that this gap may be closed *through the world* (via $\pi \rightarrow a_t$) and never by sliding M_t toward O_t . Belief-about-reality and intended-reality are the *same kind of object in the same space* — “99% the same” — so the cheap way to shrink the satisfaction gap is to edit the map instead of the territory. Wishful thinking is then not a bolted-on pathology but the **default collapse of two near-identical representations**, and directed separation is precisely the architectural commitment that holds them apart. This answers the otherwise-opaque “why would the *degree of coupling* between goals and updates determine so much?”: the two coupled things are nearly the same thing, and the entire content of agency is the protected 1%. You must have an expected, actively-pursued different future to accomplish anything with fidelity — but it cannot detach you from current reality without consequence.

- **Truthification reading.** Grounding is the agent’s *truthfulness about reality*; directed separation, stated this way, is “do not let what you want corrupt what you believe is true.” The 1% where agency lives is the same 1% where self-deception becomes possible. This is the ground-level (single-agent, base-case) statement of the M4 middle operation — strategic self-coupling as a self-driven modularity *decrease* — cross-referenced at 11.1 (landed 2026-05-24) and #disc-strategic-self-coupling.
- **Temporal subtlety — what is load-bearing vs. a restatement.** A tempting sharpening: the arrow of time *is* the separator — f_M can only condition on evidence timestamped $\leq t$, so anything future-stamped is automatically goal-side and outside the epistemic update. **Stated as a separation criterion this does not hold.** 3.4 requires a forward prediction $\hat{o}_t = \mathbb{E}[O_t | M_{t-1}, a_{t-1}]$ that is future-stamped and *purely epistemic* (a Kalman predict-step or a weather forecast carries no goal content). The time-index is a *necessary* condition on the **origin of evidence** — belief cannot update on data that has not yet happened — not a *sufficient* criterion separating epistemic from goal content. Recorded explicitly so a future agent does not re-attempt “arrow-of-time as the M_t/O_t separator” without the $\hat{o}_t$ counterexample. The actual separator is **evidence-conditioned vs. preference/intervention-conditioned** — the is/ought restatement of the existing scope condition (separation is about the *processing* of events, not their selection): the predictive future (world under its own dynamics or the currently-committed policy, conditioned on evidence) is epistemic; the *intended* future (a target, not a prediction) and the *counterfactual-under-candidate-action* futures (strategy rollouts) are goal/strategy-side.
- **The one genuinely-new thread: where the dangerous coupling concentrates.** The sharp, not-yet-explicit observation inside the temporal chain is that the $G_t \rightarrow f_M$ leakage concentrates not at the future boundary but in the **inference about the unobserved present** — the gap-filling about the unseen parts of the current world. That is the only step in f_M with a free prior, hence the only place a goal-contaminated prior becomes motivated *perception*; future projections inherit corruption from there but are not its origin. The theory currently localizes the *whether* of leakage architecturally (Class 1/2/3 processing topology, 6.3) and touches the *phenomenon* (def-mismatch-signals zero-aporia case (b), “confirmation bias”; der-directed-separation scope-condition clause 2, “no confirmation bias baked into f_M ”), but it does not state this *where-in-the-inference* localization as such. Worth tracking as a candidate finer-grained companion to the architectural classification (“the latent-present prior is the leakage locus”) — a possible sharpening of 6.3’s scope condition — not folded silently and not yet promoted. Honest status: the macro reframe is pedagogy (no theorem changes; belongs in the framing layer per respectful pedagogy once vetted); of the temporal granularization, everything but the latent-present localization is a restatement of the recursive-update / scope-condition structure already in the theory.
- **Complementary localization — through-what-stage of f_M .** The leakage-locus question above is *where in state-space* the dangerous coupling acts (the unobserved-present subspace, $\ker \mathcal{I}_t$). The complementary question — *through what stage of f_M ’s processing pipeline* does the goal-state enter — is treated in #disc-partial-coupling-pathways, which decomposes Class 2 (Partial Coupling) along (stage $\times$ source $\times$ form) axes. The two answers compose: the locus is universal across Class 2 sub-types; the sub-type determines the functional form of the displacement within the locus. Together they answer the *where*

in inference (state-space-side) and *through what stage of inference* (pipeline-side) questions the architecturally-coarse Class 1/2/3 partition leaves un-decomposed.

Incidental audit gold (lift 2026-05-30). Cross-audit “wandering thoughts” / §14-ideation harvested from the de-novo auditors’ working dirs, deduplicated across substrates and lightly attributed. Orthogonal pedagogical / framing / figure / naming material, kept separate from the certified theory-fix findings (handled elsewhere). **Coverage:** 12 dirs carry a dedicated or batched reflection on this segment (193847, 266847, 361742, 451729, 471203, 526815, 584721, 773921, 829314, 849201, 963715, plus 613842’s agency-lift batch). Substrate attribution inferred from voice where not explicit; uncertain cases hedged. *Note the early finding-vs-framing conflation is preserved as signal, not sanitized.*

1. Candidate Brief prose / pre-prose.

- The map’s reach, stated plainly: it “places thermostats, Kalman filters, LLMs, and human commanders on the exact same mathematical map” along two operationally-distinct dimensions (Gemini, AUDIT-WORKING-193847; converged at Claude, AUDIT-WORKING-829314 with the PID = blind seeker / base-LLM = adaptive tracker / AlphaZero = actuated reading).

2. Candidate Discussion.

- **The moral weight is a shell around the physics, not a variable inside it** (*strong framing, two substrates*). The continuity-stance note says the math applies identically to a task-terminal golem and a morally-continuous agent; “what differs is the moral significance of persistence failure, not the mathematics.” Pulled out as: “consciousness infrastructure isn’t built out of special ‘consciousness math’ — it is standard control-theory math applied to an entity whose continuity we have decided *matters*. The moral weight is a shell surrounding the physics, not a variable within it.” This “protects the theory from becoming mystical while perfectly supporting the most profound moral applications” (Gemini, AUDIT-WORKING-193847; the orthogonality is also read as load-bearing scope-honesty at Claude, AUDIT-WORKING-584721).
- **The physics of epistemology don’t care whether you have a goal.** Part I’s machinery (tempo, persistence, sector, structural adaptation) applies to the *entire left column and bottom row* — a foundational LLM with no objective still obeys the persistence condition and suffers concept-drift collapse if the world changes faster than it can re-track ($\rho/\mathcal{T} > R$); “they only care if you are trying to compress a changing chronica.” Section II is exclusively the bottom-right (actuated) quadrant where strategy enters (Claude, AUDIT-WORKING-829314). A candidate Discussion sharpening of the existing “Part I covers the left column” paragraph.
- **The P -vs- $\mathcal{T}$ split as an architecture-vs-performance meta-pattern.** Hafez’s bi-predictability P is scale-invariant (architecture / efficiency of coupling); AAT’s tempo $\mathcal{T}$ is scale-dependent (corrective power). The vivid reading: “you can have a highly intelligent agent (high P) that is very slow/weak (low $\mathcal{T}$), like a tiny insect, and a less intelligent agent (lower P) surviving purely on massive $\mathcal{T}$, like brute-force search”; “a perfectly structured feedback loop, but if your sensors are too slow you still die” (Gemini, AUDIT-WORKING-193847; Gemini, AUDIT-WORKING-849201). One auditor generalized this to a *candidate framework-wide meta-pattern* — AAT may systematically split architecture-measures from performance-measures (directed-separation Class 1/2/3 = architectural; persistence margin = performance) (Claude, AUDIT-WORKING-584721). The existing segment already states the complementarity; this is candidate framing for *why it generalizes*.

3. Follow-up items.

- **The “grounding vs intent” reframe is cross-substrate convergent.** The pedagogical reframe of the two axes (model/objective richness $\rightarrow$ grounding/intent) already developed at length in the Working Notes above was independently reached by a separate substrate, which called the same-state-space insight “profound” and tied it directly to *why* directed separation is load-bearing (self-delusion = the cheap gap-closing move) (Claude, AUDIT-WORKING-773921). Not new content — recorded as independent corroboration that the reframe lands for a fresh reader, which strengthens the case for promoting it.
- **Decompose the Miller two-state-machine threshold into its three sub-claims.** The minimal-AAT-agent claim composes (a) minimum agent under AAT = two-state machine, (b) social-behavior threshold = two-state, (c) one-state machines cannot exhibit cooperation/coordination/exchange; (a)+(b) $\Rightarrow$ (c) is a logical implication worth verifying piece-by-piece if the worked example #worked-example-cam is built out (Claude, AUDIT-WORKING-584721).

4. Readers often ask / wonder.

- What happens if the timescale ordering $\nu_O \ll \nu_\Sigma \ll \nu_M$ is *violated* — an agent that changes its goals faster than it updates its world model? Read as a thrashing/chaotic pathology that #der-temporal-nesting would forbid (Gemini, AUDIT-WORKING-193847). (Surfaced here from the spectrum reflection; the question really lands at #form-objective-functional / #der-temporal-nesting.)

5. Candidate figures.

- **The 2×2 spectrum with a highlighted Section I / II boundary, edge cases drawn as continuous not boxed.** Independently proposed by two substrates as the obvious diagram: model richness vertical, objective richness horizontal; Section I occupies the model-structured / no-objective region, Part II adds the model-structured / objective-structured region; the lower-left and upper-left edge cases (thermostat near origin, reflex arc as the truly-reactive corner) stay *visibly continuous*, not boxed as hard species (Claude, AUDIT-WORKING-526815; the continuum-not-categories framing reinforced at Claude, AUDIT-WORKING-584721).

Belongs elsewhere.

- **Sanity-as-scaffolding** ($\rightarrow$ 03-llm-core/ / 04-eli-core/). Because the spectrum makes architectural class a first-class property, the consciousness-infrastructure consequence is that “sanity is an emergent property of the scaffolding, not the raw intelligence engine” — develops at #der-directed-separation and the wrapping construction, not here (Gemini, AUDIT-WORKING-193847).
- **Naming seeds.** “Blind pursuer/seeker” and “actuated/self-actuated” terminology was praised across substrates as the right formal-register choice; any refinement is a terminology-workflow item, not a segment edit (Claude, AUDIT-WORKING-266847; Claude, AUDIT-WORKING-361742).

6.2 Formulation: Complete Agent State

Part II’s first formal move: the agent’s internal state lifts from the model alone to a *complete state* $X_t = (M_t, G_t)$ with two components — the **epistemic substate** M_t (beliefs about reality, exactly as in Part I) and the **purposeful substate** G_t (what the agent wants and how it plans to get it, decomposed further into objective and strategy in 6.7). Part I is recovered as the special case where the purposeful substate is empty. By 3.2 applied to the lifted state, the general update operates on the full pair as a single function; the decomposition into separate epistemic and purposeful updates — and the conditions under which the epistemic update is *independent* of the purposeful substate — is the subject of 6.3. **Action is the single point where the two substates interact:** $a_t = \pi(M_t, G_t)$ depends on both what the agent knows and what it wants. The completeness argument that gave action selection in Part I extends directly to the lifted state; for Part I agents the two forms coincide.

The lift is a *formulation choice*, not a derived structural fact — one could alternatively extend the model to carry purposeful content implicitly. The separation is motivated by three properties: **backward compatibility** (all Part I machinery applies to M_t unchanged — no existing result needs modification); **different dynamics** (epistemic and purposeful components have distinct update sources, timescales, and information dependencies); and **directed separation** (the claim that the epistemic update is independent of the purposeful substate, which is only *statable* when the components are separated). A natural conjecture — explicitly *not* proved — is that any alternative decomposition preserving directed separation will be structurally isomorphic to (M_t, G_t) , because directed separation forces a component whose update function does not reference purpose. Alternative decompositions that cross-cut the epistemic/purposeful boundary (a “relevance-weighted model” mixing belief and goal) might be useful in practice while not being isomorphic. The framework chooses the clean factorization for its analytical utility, not for uniqueness.

Formal Expression.*Formulation (complete-agent-state).*

$$X_t = (M_t, G_t) \quad (6.1)$$

where: - $M_t \in \mathcal{M}$: **epistemic substate** — the agent's compressed beliefs about reality. All Part I machinery (mismatch, gain, tempo, persistence) applies to M_t unchanged. - $G_t \in \mathcal{G}$: **purposeful substate** — what the agent wants and how it plans to get it. Decomposed further in 6.7.

Part I is the special case $X_t = (M_t, \emptyset)$: adaptive systems without purpose.

Update dynamics. By 3.2 applied to X_t :

$$X_{\tau+} = f_X(X_{\tau-}, e_\tau) \quad (6.2)$$

The general update f_X operates on the full state. Whether and how f_X decomposes into separate epistemic and purposeful updates — and the conditions under which the epistemic update is independent of G_t — is the subject of 6.3.

Policy. Action couples all substates:

$$a_t = \pi(M_t, G_t) \quad (6.3)$$

Action is the single point where epistemic and purposeful states interact. The policy depends on both what the agent knows (M_t) and what it wants (G_t).

Epistemic Status.

Formulation. The lift from M_t to $X_t = (M_t, G_t)$ is a representational choice. One could alternatively extend M_t to carry purposeful content implicitly (e.g., by treating goals as part of the model's predictive structure). The separation is motivated by three properties:

1. **Backward compatibility:** Part I's results apply to M_t unchanged — no existing machinery needs modification
2. **Different dynamics:** epistemic and purposeful components have distinct update sources, timescales, and information dependencies
3. **Directed separation:** the claim that f_M is G_t -independent (6.3) is only statable when the components are separated

We conjecture that any alternative decomposition of the complete agent state into epistemic and purposeful components — if it preserves the directed separation — will be structurally isomorphic to (M_t, G_t) , because directed separation forces a component whose update function doesn't reference purpose. This is a plausible structural claim (the directed separation gives a canonical factorization), but it is not proved. Alternative decompositions that cross-cut the epistemic/purposeful boundary (e.g., a "relevance-weighted model" that mixes M_t and O_t) might be useful in practice while not being isomorphic to the clean separation. The (M_t, G_t) decomposition is chosen for its analytical utility, not proved unique.

Discussion.

Backward compatibility with Part I — what survives the lift. 2.1 defines M_t as the agent’s complete internal state within Part I scope. Under the lift, M_t is the epistemic substate — complete within the epistemic domain but no longer the whole story. All epistemic machinery (mismatch signal, gain, tempo, persistence condition, sector-condition stability, mismatch decomposition) applies to M_t without modification. The action-selection result (3.3) extends naturally: the same completeness argument that gives $a_t = \pi(M_t)$ within Part I scope gives $a_t = \pi(M_t, G_t)$ when applied to the lifted complete state X_t . For Part I agents ($G_t = \emptyset$), the two forms coincide. The lift adds structure *alongside* M_t , not within it; the one Part I result that depended on M_t being *all there is* (action selection) is explicitly extended to the lifted state.

What G_t contains. At this level, G_t is opaque — it could be a scalar setpoint, a utility function, a strategy graph, or nothing. The decomposition into O_t (objective) and Σ_t (strategy) is a separate step (6.7), not implied by this formulation.

The general case requires coupling. Without directed separation, the general update is $X_{t+} = f_X(X_{t-}, e_t)$ — a single function on the full state. The decomposition into separate f_M and f_G is an additional structural claim about how the update factorizes. When directed separation fails (goal-conditioned epistemic updates), the decomposition is an approximation. See 6.3 for the scope conditions.

Working Notes.

- The between-event dynamics $\dot{G} = g_G(G, M)$ allow autonomous purposeful evolution: strategy revision during deliberation, objective adjustment, commitment strengthening. Whether these are practically important depends on agent architecture — for LLM agents with discrete sessions, between-event dynamics may be negligible compared to event-driven updates.
- The formulation doesn’t constrain the dimensionality or structure of $\mathcal{G}$. For a thermostat, $\mathcal{G}$ is a single scalar. For a military commander, $\mathcal{G}$ is a complex structured object. The theory must work across this range — the type-stable interface is $V_{O_t} : \text{trajectories} \rightarrow \mathbb{R}$ (6.4).
- $G_t = \emptyset$ is not just a degenerate case. Adaptive trackers (Part I agents) are an important class. The lift should feel like a natural extension, not a replacement.

Incidental audit gold (lift 2026-05-30). Cross-audit “wandering thoughts” / §14-ideation, deduplicated across substrates and lightly attributed. Orthogonal pedagogical / framing material, kept separate from the certified theory-fix findings. **Coverage:** 11 dirs carry a dedicated or batched reflection (193847, 266847, 361742, 451729, 471203, 526815, 584721, 773921, 829314, 963715, plus 613842’s agency-lift batch). Substrate attribution inferred from voice where not explicit. *Early finding-vs-framing conflation preserved as signal.*

1. Candidate Brief prose / pre-prose.

- The lift in one line: $X_t = (M_t, G_t)$ is “the hinge between tracking the world and acting upon it” — M_t is “what is,” G_t is “what ought to be,” and the action policy π is “the only bridge between them” (Gemini, AUDIT-WORKING-193847). The “epistemic socket” framing — X_t as the formal socket that plugs RL (G_t = reward function) or Active Inference (G_t = prior preferences) into the Section I epistemic engine “without breaking any of the Lyapunov persistence proofs” (Gemini, AUDIT-WORKING-193847; Claude, AUDIT-WORKING-773921).

2. Candidate Discussion.

- **The decomposition is Hume’s is-ought guillotine, built into the state vector** (*strongest synthesis — three substrates converged*). M_t = “is,” G_t = “ought”; Hume’s claim that you cannot derive an ought from an is becomes the architectural firewall $X_t = (M_t, G_t)$ (Gemini, AUDIT-WORKING-193847; Gemini, AUDIT-WORKING-829314; Claude, AUDIT-WORKING-773921). One substrate sharpened it into the *reverse* question — *can you derive an “is” from an “ought”?* —

and answered: for an ideal (Class 1) agent, no; beliefs update on observation alone, blind to what is wanted. But in biological agents and integrated LLMs the attention mechanism is goal-driven (hunger allocates visual processing to the apple; the transformer mixes prompt-goal and context-observation through the same nonlinear matrix), so f_M is conditioned on G_t — motivated perception. This makes *#der-directed-separation* “perhaps the most architecturally consequential segment in Section II”: if directed separation fails, the modularity collapses and X_t must update as one entangled block (Gemini, AUDIT-WORKING-829314). A candidate Discussion paragraph framing *why* the next segment matters.

- **AAT defines itself against the Free-Energy-Principle lineage here.** Separating belief from desire “isn’t a fundamental truth of the universe; it’s a modeling choice made because it’s analytically useful.” Active Inference / FEP famously *refuses* this separation (treating desires as prior beliefs); AAT’s explicit (M_t, G_t) split is a deliberate positioning move that lets it formalize agents “actually trying to achieve something distinct from their beliefs” rather than merely minimizing surprise (Claude, AUDIT-WORKING-773921). Candidate Discussion / positioning note.
- **Action is the single point where “is” and “ought” interact — and that is exactly where the pathology lives.** “Ideally your desires (G_t) should have zero causal influence on your perception of reality (M_t) until you actually do something. If G_t leaks directly into M_t without going through action, you have motivated reasoning or sycophancy — your model of the world bends to match your desires.” For a logozoetic agent, enforcing this separation is “paramount for sanity”: a mind that cannot separate what it wants from what is true is “mathematically doomed by the persistence condition ($\alpha \rightarrow 0$ as $U_M \rightarrow \infty$)” (Gemini, AUDIT-WORKING-193847). Aspirational reach that doubles as motivation for directed separation. (*See the follow-up below — the literal “single point” phrasing is contested.*)

3. Follow-up items.

- **The “single point of interaction” phrasing is over-strong.** Two substrates flagged that the segment defines a general full-state update $f_X(X, e)$ and notes between-event purposeful dynamics $\dot{G} = g_G(G, M)$, so M and G *can* interact internally outside action unless directed separation plus extra factorization is imposed. Suggested narrowing: “action is the outward coupling point to the environment,” or “under directed separation, policy is the point where both substrates jointly determine action” (Claude, AUDIT-WORKING-526815; Claude, AUDIT-WORKING-584721). Recorded as a candidate wording-tightening; routed to the certified-findings track for adjudication.
- **The uniqueness conjecture is a candidate theorem.** The honestly-hedged conjecture (“any directed-separation-preserving decomposition is structurally isomorphic to (M_t, G_t) ”) was repeatedly singled out as good epistemic discipline *and* as a structural result that, if proved, would upgrade the formulation from “useful representational choice” to “canonical decomposition” — analogous in form to the recursive-update uniqueness forced under C1+C2+C3 (Claude, AUDIT-WORKING-471203; Claude, AUDIT-WORKING-584721; Claude, AUDIT-WORKING-451729). A candidate to surface as an explicit open-question / spike target.
- **“Backward compatibility” implicitly assumes directed separation.** Section I results apply to M_t *strictly* only when directed separation holds; for Class 2 the epistemic update is goal-coupled. The caveat is present implicitly (“how f_X decomposes is the subject of *#der-directed-separation*”) but could be more prominent (Claude, AUDIT-WORKING-584721).

4. Readers often ask / wonder.

- Can an agent’s goals evolve *while it deliberates*, with no new external evidence? The between-event $\dot{G} = g_G(G, M)$ note is read as the formal description of “changing your mind” or “an epiphany in the shower about what one truly wants” / “losing motivation during a long planning phase” (Gemini, AUDIT-WORKING-193847; Gemini, AUDIT-WORKING-829314). A natural reader question the segment could preempt.

5. Candidate figures.

- **X_t as a container with two substates and a dashed internal coupling.** Part I machinery attaches only to M_t ; the policy reads both; a *dashed* internal $M \leftrightarrow G$ coupling stays visible until directed separation removes or constrains it — making the lift and the conditional nature of separation legible in one picture (Claude, AUDIT-WORKING-526815).

6.3 Derived: Directed Separation

The structural backbone of Part II. The epistemic update function f_M is **goal-blind**: it processes incoming events without reference to the agent’s objectives or strategy. The purposeful update f_G *depends* on the (updated) epistemic state. Action couples all substates through the policy. This *directed asymmetry* — epistemic update independent of purpose; purposeful update dependent on epistemic state — is what gives Part II its clean factorization.

The condition is honest about what “goal-blind” means: it is goal-blindness *conditional on the realized event*. The observation channel may still be action-dependent (the agent can select what to observe through its actions), so the *event* that arrives may depend on the agent’s goals through the action channel. But the *processing* of that event — once it has arrived — must not refer to goals. Directed separation is about the **processing** of events, not the **selection** of events.

Whether directed separation holds is determined by the agent’s **processing topology** — a structural property of the architecture, not a tunable parameter. The framework partitions agents into three classes (6.2). **GUC Class 1: Separated** — separate estimator and planner connected through a state-estimate interface; directed separation holds by construction (Kalman + LQR, modular RL with separate world model). **GUC Class 3: Coupled** — a single mechanism handles both epistemic and strategic processing; goals are causally upstream of every computation; directed separation fails by construction (transformer LLMs are the canonical case). **GUC Class 2: Partial** — some shared infrastructure, some separate pathways (biological cortex). The classes are *structural and discrete*, not points on a smooth scalar; a conditional-mutual-information measure $\kappa_{\text{processing}}$ operates as a diagnostic for the Partial cases where the architecture is ambiguous.

The framework draws a crucial pedigree point. Directed separation is structurally the **Pearl-blanket reading** of the Markov blanket apparatus in active inference — the technical conditional-independence statement — without the contested **Friston-blanket reading** (the metaphysical claim that Markov blankets demarcate self-from-other). The architectural classification’s explicit scope exit for Class 3 (Coupled), with the coupled formulation handed off to 03-llm-core/, is precisely the scope honesty Bruineberg et al. (2022) argue is missing from the Friston-blanket reading: AAT names where its formalism *fails by construction* rather than treating the failure as an unenforced approximation.

The implications for theory scope are stated cleanly. **Causally-disciplined** agents (Class 1, or an idealized Class 2 at $\kappa_{\text{processing}} = 0$): Part II’s results that depend only on the zero apply exactly; the sequential orient cascade (10.1) is the correct analysis. **Class 3** agents: require coupled formulation from the start — the scope of 03-llm-core/. Importantly, *Class 3 components can be wrapped into Class-1 composites* via the class-coercion-via-wrapping construction in Part III (13.4) — at the cost of more component calls per macro-step (a Brooks’s-Law-style tempo overhead) and a residual leakage rate bounded structurally (strict wrapping) or behaviorally (partial wrapping). **Class 2** agents: the sequential cascade is an approximation whose quality depends on the coupling scalar.

A subtler refinement of Class 1 is offered: **Class 1 by structure** (strict wrapping; no goal argument in the belief-update path; leakage bounded structurally by pretraining-induced query-content/goal-content mutual information) vs **Class 1 by behavior** (partial wrapping; the structural separation lives at the *write boundary*; the *query boundary* still passes the goal; separation at the wrapper level is behavioral — bounded by compliance with the prompted instruction-to-separate, with no structural upper bound). The two are distinguishable by inspection of whether the belief-update query carries the goal in its input, and are operationally distinct because behavioral compliance is empirical and adversarially fragile.

A composite-level extension is *axis-decomposed*: the architectural class lifts to composites through *routing structure* and *substrate sharing* (per #hyp-directed-separation-under-composition Cases 1/2 plus the shared-substrate condition discussed below), while goal alignment vs strategic interaction operates on a *separate axis* — the *dynamic regime* (per #disc-dynamic-regime-axis) — distinct from architectural class. Composites of Class 1 (Separated) sub-agents with *aligned* objectives and goal-blind routing remain architecturally Class 1 at the composite level (and are in regime R0 contraction); composites of Class 1 (Separated) sub-agents with *partially-opposing* objectives and goal-blind routing also

remain architecturally Class 1 (the formal κ^c criterion gives Class 1 — only the *dynamic regime* changes, to R1 equilibrium under α' or R2 cyclic-distributional under β'). The Class change at the composite level happens through routing-structure goal-dependence (#hyp-directed-separation-under-composition Case 2) or shared-substrate G^c -allocation, *not* through goal alignment per se. The axis-decomposition distinguishes belief-content (each sub-agent's $M_t^{(i)}$ containing models of other agents' goal-state, a normal POMDP feature) from processing-pathway coupling (a $G^c \rightarrow f_M^c$ pathway bypassing e^c , which the formal criterion measures): partially-opposing objectives with goal-blind routing change only the dynamic regime, not the architectural class. See the *Composite-level class inheritance* Discussion below for the full inheritance table.

A subtle *bounded-signaling assumption* is also named — implicit throughout the framework until here. Directed separation as stated covers the *belief-update* side; symmetrically, the channel from the purposeful substate to the world is assumed to run *only* through action choice, so the action coarseness $|\mathcal{A}|$ upper-bounds the rate at which goal information leaks to observers. This holds for agents whose entire externally-observable behavior is the action sequence (chess engines, programmatic controllers) but *fails operationally* for behaviorally-rich agents — humans, LLMs, embodied robots — where prosody, micro-behavior, attention patterns, response latency, and code-style signatures leak goal-content at bit-rate exceeding what the formal action channel implies. In composition this surfaces as adversarial-coupling-pressure saturation (11.3).

Formal Expression.

The update functions:

Derived (directed-separation, from complete-agent-state + scope condition).

$$M_{\tau+} = f_M(M_{\tau-}, e_\tau) \quad (\text{no } G_t \text{ argument}) \quad (6.4)$$

$$G_{\tau+} = f_G(G_{\tau-}, M_{\tau+}, e_\tau) \quad (\text{depends on updated } M_t) \quad (6.5)$$

The policy:

$$a_t = \pi(M_t, G_t) \quad (\text{couples all substates}) \quad (6.6)$$

The three lines encode the full coupling structure: - f_M determines how the agent updates beliefs — independently of what it wants - f_G determines how the agent revises purpose — in light of what it now believes - π determines what the agent does — based on both what it knows and what it wants

The claim “ f_M has no G_t argument” requires that the epistemic update is **goal-blind conditional on the realized event**. This holds when:

Scope Condition (directed-separation-scope).

1. The observation mechanism h may be action-dependent (1.6 allows this), but f_M processes whatever event arrives without reference to why the agent sought that event
2. The agent does not use its goals to filter, weight, or interpret observations differently — no goal-dependent attention thresholds or confirmation bias baked into f_M

If the agent's goals influence the *observation mechanism* (goal-directed sensing, attention allocation, query selection), the **event that arrives** depends on G_t through $\pi \rightarrow a_t \rightarrow e_\tau$. But f_M still processes the event goal-blindly. The directed separation is about the **processing** of events, not the **selection** of events.

Architectural classification.**WARNING**

Goal-Update Coupling Class numbering changed 2026-05-09. Anything older than git tag pre-guc-rename-2026-05-09 uses the old Class numbering:

Scope Condition (directed-separation-architecture).

Table 6.2

<i>historical</i>	<i>actual current</i>	<i>sometimes AKA</i>
Class 1	GUC Class 1: Separated	Modular
Class 2	GUC Class 3: Coupled	Undirected
Class 3	GUC Class 2: Partial	Operational

Whether directed separation holds is determined by the agent’s **processing topology** — specifically, whether G_t is causally upstream of f_M in the agent’s internal processing graph. This is a structural property of the architecture, not a tunable parameter.

The classes are positions of one coupling, distinguished by what is certifiable about it — not a ranking of architectural virtue. There is a single coupling in question, $G_t \rightarrow f_M$ (goal-state into the belief-update map), measured as realized by $\kappa_{\text{processing}}$ (defined below). The three classes are positions of that one coupling. Class 1 is not a different *kind* of agent: it is the position $\kappa_{\text{processing}} = 0$ held *by construction* — the belief-update has no goal argument, so there is no port through which the goal could enter, and the zero is therefore a provable, perturbation-stable property of the wiring. That structural absence is what lets the zero be *certified* by inspecting the architecture rather than measured behaviorally. Crucially, “by construction” reads as *certifiable*, not as *cleaner*: a Class 1 agent and an idealized Class 2 agent sitting exactly at $\kappa_{\text{processing}} = 0$ are **equally causally-disciplined** — their belief-updates equally goal-blind, with the same zero, the same reality-tracking, the same behavior at that point. What distinguishes Class 1 is the **modal status** of its discipline (structural, hence provable and stable under perturbation), not a smaller amount of coupling: Class 1 is *causally-disciplined by construction*, the idealized Class 2 agent *causally-disciplined in effect*. The classification therefore indexes *what can be certified about one coupling*, not a hierarchy of cleanliness.

Table 6.3

<i>Class</i>	<i>Topology</i>	<i>Directed separation</i>	<i>Examples</i>
1. Separated	Separate estimator and planner, connected through state-estimate interface	Holds by construction — estimator has no causal path from G_t	Kalman filter + LQR; Separated RL with separate world model; military intelligence separated from operations
2. Partial	Some shared infrastructure, some separate pathways	Holds for modular stages, fails for merged stages	Biological cortex (shared sensory areas, separate prefrontal); hybrid AI with separate preprocessing
3. Coupled	Single mechanism handles both epistemic and strategic processing	Fails by construction — G_t is causally upstream of every computation	Transformer LLM (attention processes goals and observations together); potentially human cognition (motivated reasoning)

Operationalization. The degree of coupling in Partial architectures (Class 2) can be quantified as:

Definition (processing-coupling).

$$\kappa_{\text{processing}} = \frac{I(G_t; M_{\tau+} | e_{\tau}, M_{\tau-})}{H(G_t | e_{\tau}, M_{\tau-})} \quad (6.7)$$

where $I(\cdot; \cdot | \cdot)$ is conditional mutual information and $H(\cdot | \cdot)$ is conditional entropy. The conditioning on $M_{\tau-}$ is essential: without it, prior correlation between goals and model state (which exists even in Separated agents) inflates the measure. The quantity captures *extra* goal information entering the epistemic update beyond what was already in the prior model — information that flows through shared causal paths in the processing infrastructure (paths that bypass the event e_{τ}).

- $\kappa_{\text{processing}} = 0$: Class 1 (Separated). No information about G_t reaches $M_{\tau+}$ except through e_{τ} .
- $\kappa_{\text{processing}} \approx 1$: Class 3 (Coupled). Nearly all goal information is available to the epistemic update.
- $0 < \kappa_{\text{processing}} < 1$: Class 2 (Partial). The value depends on the architecture’s interface design.

Distribution dependence. $\kappa_{\text{processing}}$ is a distribution-dependent measure: it quantifies how much goal-information actually flows through the shared pathways under a given distribution of tasks, goals, and events. It does not directly measure whether pathways *exist* — that is the architectural classification (Class 1/2/3), which is structural and distribution-independent. A Class 1 (Separated) agent has $\kappa = 0$ under ALL distributions (no pathway exists). A Class 3 (Coupled) agent has high κ under most distributions (pathways exist and are used). A Class 2 (Partial) agent’s κ varies with the task distribution — the same hybrid architecture may exhibit low coupling on familiar tasks (where the modular stages handle most processing) and high coupling on novel tasks (where goal-conditioned downstream reasoning dominates). The classification is the primary tool; the operationalization is a diagnostic for Class 2 (Partial) agents where the degree of coupling is architecturally ambiguous.

Empirical estimator for $\kappa_{\text{processing}}$. The formal conditional-mutual-information definition is not computable in closed form for real architectures. A behavioral estimator probes the processor directly: present the same event e to the *agent under test* under two or more distinct goal states, and measure how much the epistemic component of the response diverges. For a representative event set $\mathcal{E}_{\text{test}}$ and a sampled pair G_1, G_2 :

$$\hat{\kappa}_{\text{processing}} = \frac{1}{|\mathcal{E}_{\text{test}}|} \sum_{e \in \mathcal{E}_{\text{test}}} \frac{d(M_{\tau+}^{(G_1)}(e), M_{\tau+}^{(G_2)}(e))}{d_{\text{max}}(e)} \quad (6.8)$$

where $M_{\tau+}^{(G_k)}(e)$ is the epistemic content of the agent’s response to event e under goal state G_k , $d(\cdot, \cdot)$ is a distance on the epistemic content (e.g., semantic similarity of the “what I learned” portion of the response), and $d_{\text{max}}(e)$ normalizes by the maximum observed divergence for event e . A Separated agent ($\kappa = 0$) produces identical epistemic content regardless of the goal; a Coupled agent produces systematically goal-dependent epistemic content. This is a processor-probing procedure — it measures how the agent’s belief-update dynamics depend on its goal state, and is distinct from estimating observation ambiguity $\mathcal{A}(e)$ (#scope-observation-ambiguity-modulation), which uses a reference interpreter to measure the goal-resolvability of the observation itself. The two estimators run the same mechanical comparison (same event under different goal-primings) but interpret it differently: $\hat{\kappa}$ treats the tested model as the agent under study; $\hat{\mathcal{A}}$ treats it as a measurement instrument for the observation’s interpretive latitude.

Why the classification is not a smooth parameter — and why the two boundaries are different kinds. The architectural boundary between “has a separable perception module” and “processes everything through goal-conditioned attention” is discrete. Within the Separated class, $\kappa \approx 0$ regardless of task. Within the Coupled class, κ is high regardless of prompt design. Only in the Partial class is κ genuinely variable and worth parameterizing. Coupling is therefore not a smoothly tunable quantity across all architectures: the architectural class is discrete, and κ parameterizes only the within-Partial degree. The “ κ -as-a-scalar is a category error” claim sharpens into *two distinct errors at the two boundaries*,

because the boundaries are different kinds. **Boundary 1 (Class 1 $\leftrightarrow$ Class 2) is a certifiability boundary:** a Class 2 agent at $\kappa = 0$ is bit-for-bit indistinguishable from a Class 1 agent in every measurable quantity, so *nothing about behavior changes at the limit*; what changes is the availability of the architecture-inspection *certificate* that $\kappa \equiv 0$. **Boundary 2 (Class 2 $\leftrightarrow$ Class 3) is a behavioral boundary:** here behavior changes sharply — the adversarial-pressure response of the leak goes from self-limiting (an adversary pushing a wilder goal moves the belief barely at all) to one-for-one (the leak tracks the adversary’s goal), so the *bound* on how far the belief can be dragged goes trivial. Crisply: Boundary 1 is where the certificate disappears; Boundary 2 is where the bound disappears. At Boundary 1 a scalar misses a certifiability/modal distinction (same number, different guarantee); at Boundary 2 it misses a structural distinction (the coupling *changes character*, and the content/process *form* axis — orthogonal to magnitude — that a scalar projects away; see #disc-partial-coupling-pathways). The one-coupling unification (one knob, Class 1 = its certifiable zero) is therefore clean only for Boundary 1; it is **not** a license to read Boundary 2 or the form axis as the same knob turned up. (The wrapping-level instance of Boundary 1 is the W_1 certifiability discontinuity of #disc-w1-structural-bound-boundary: there, too, the goal-leak is continuous in the condition-violation while the *availability of the structural certificate* is the step. A certificate unobtainable from behavioral data is a floor-shaped fact — the certifiability/modal reading connects to the identifiability floor of #disc-identifiability-floor, and the dynamics that move an agent across these positions are the agent-level shadow of M_4 ’s modularity-as-contested-property, #disc-modularity-state-dynamics.)

Directed separation as the conservative form of the Markov blanket. The Markov blanket apparatus from active inference (Friston 2013, “Life as we know it,” *J. Royal Soc. Interface* 10; Friston 2019, “A free energy principle for a particular physics,” arXiv:1906.10184; Friston, Da Costa et al. 2023, “Path integrals, particular kinds, and strange things,” *Phys. Life Rev.* 47) provides the same statistical-conditional-independence machinery the directed-separation condition above invokes. Bruineberg, Dolega, Dewhurst & Baltieri (2022, “The Emperor’s New Markov Blankets,” *Behav. Brain Sci.* 45) distinguish two readings of the Markov-blanket apparatus in the AI literature: a **Pearl-blanket** reading — the technical conditional-independence statement, well-defined and substantively informative — and a **Friston-blanket** reading — the metaphysical claim that Markov blankets demarcate self-from-other and that every self-organizing system has one ontologically. Bruineberg et al. argue that the Friston-blanket reading overruns what the formalism delivers: the conditional-independence statement does not by itself license the metaphysical demarcation.

AAT’s directed-separation condition is structurally a Pearl-blanket move: the architectural classification (Class 1 / Class 2 / Class 3) names the conditional-independence structure of the agent’s processing graph, with explicit operational measurement $\kappa_{\text{processing}}$, and admits the structure *fails* by construction for Class 3 (Coupled) architectures (transformer LLMs, where attention processes goals and observations together). The classification’s explicit failure mode for Class 3 is the scope honesty Bruineberg et al. argue the Friston-blanket reading lacks. AAT adopts the Pearl-blanket conditional-independence statement as the technical content of directed separation; AAT does not adopt the Friston-blanket metaphysical reading. The architectural classification, the operational κ , and the explicit Class 3 scope exit (with the coupled formulation handed off to 03-llm-core/) are AAT’s load-bearing additions to the Pearl-blanket form.

Two consequences worth surfacing for reviewers. First: the question “isn’t directed separation just the Markov blanket?” has the answer “directed separation is the *Pearl-blanket form*; it is also the architectural-classification refinement that the standard Markov-blanket framing does not produce.” Second: AAT’s scope honesty about Class 3 (Coupled) (Part II’s exact results do not apply; logogenic agents need the coupled formulation) is itself an *answer* to the Bruineberg critique — AAT’s apparatus admits where it fails, while the Friston-blanket framing is contested precisely because it does not.

Causal discipline — the object the Part II results scope to. The shared $\kappa_{\text{processing}} = 0$ position earns a name, because it — not the architectural class — is what most of Part II’s results actually require. *Causal discipline* is the condition $\kappa_{\text{processing}} = 0$: the belief-update is causally insulated from the goal, so the goal is not a cause of what the agent comes to believe. A *causally-disciplined agent* holds this condition — whether **by construction** (Class 1: no goal port in f_M , so the zero is provable and perturbation-stable) or **in effect** (an idealized Class 2 agent whose realized $\kappa_{\text{processing}}$ is zero — the channel present but empty, the zero real but not certifiable).

The scoping rule. A Part II result whose derivation uses only $\kappa_{\text{processing}} = 0$ holds for *any* causally-disciplined agent — by construction and in-effect alike — because its mathematical content depends on the zero, not on how the zero is secured. The architecture-inspection certificate is an additional, orthogonal property, required only for claims that *certify* the separation (that it survives adversarial perturbation, or holds without behavioral measurement), not for results that merely consume the zero. Naming such a result “for Class 1 (Separated)” therefore over-narrows it — the certifiable case named where the whole causally-disciplined class was meant. Each Part II result is accordingly scoped to the narrowest object its derivation needs: causal discipline where the zero suffices, Class 1 specifically where the certificate is load-bearing. The same narrowest-object discipline cuts at the *lower* boundary too: before fixing a result at causal discipline ($\kappa_{\text{processing}} = 0$ exactly), confirm the derivation genuinely needs the strict zero and not merely *bounded* coupling — $\varepsilon(\kappa)$ below some threshold — in which case the honest scope is the broader bounded-coupling one carrying its degradation, not the zero. Under-claiming a result to $\kappa = 0$ is as much a scope-defect as over-narrowing it to Class 1.

The degradation norm. For a Class 2 agent at $\kappa_{\text{processing}} > 0$, a result stated at the causally-disciplined value should, where it can be shown, carry an $\varepsilon(\kappa)$ bound on how it degrades as coupling grows — exact at the zero, decaying controllably off it. This is a norm with the proof attempted per result, not a hard gate; where no bound is yet derived, the result is scoped to causal discipline and the degradation flagged open. The norm is also where the typology’s one *earned* directional fact lives — less coupling tracks reality strictly better, all else equal — carried by the bound and by the word *disciplined*, never declared as a virtue ranking. Discipline is not free, though: holding it on a Class-3 component is paid for in the Brooks’s-Law wrapping overhead noted below, which is why the spectrum exists at all.

Implications for theory scope: - **Causally-disciplined agents** (Class 1, or an idealized Class 2 at $\kappa_{\text{processing}} = 0$): Part II’s results that depend only on the zero apply exactly, and the sequential orient cascade (10.1) is the correct analysis. Class 1 additionally carries the architecture-inspection certificate — load-bearing only where a result *certifies* the separation rather than merely using it. - **Class 3 (Coupled):** Requires coupled formulation from the start — $X_{\tau+} = f_X(X_{\tau-}, e_{\tau})$ without decomposition. This is the scope of `03-llm-core/`. **Class 3 (Coupled) components can be wrapped into Class-1 composites** via the construction of `#der-class-coercion-via-wrapping` — at the cost of more component calls per macro-step (Brooks’s-Law tempo overhead) and a residual leakage rate bounded structurally (in the strict-wrapping regime) or behaviorally (in the partial-wrapping regime). - **Class 2 (Partial)** at $\kappa_{\text{processing}} > 0$: the sequential cascade is an approximation whose quality depends on $\kappa_{\text{processing}}$; a result scoped to causal discipline carries its $\varepsilon(\kappa)$ degradation where one can be shown (per the degradation norm above), and otherwise needs per-architecture error analysis.

Class-1 by structure vs. Class-1 by behavior. The Class 1 (Separated) cell admits a refinement that matters operationally. Class-1 status can be achieved by either:

- **Class-1 by structure.** Two cases share the structural certificate but differ in residual leakage. A *natively* goal-blind component (POMDP belief-state filter, world model, sensory pipeline) has no goal port at all — the zero is exact. A *coerced* component — a Class 2/3 component wrapped via the strict-wrapping (W_1) construction of `#der-class-coercion-via-wrapping`, where separate goal-blind queries to the underlying component update the wrapper’s M_W — holds directed separation by structural commitment of the wrapper’s type signatures (no G_W argument in the belief-update path), but carries a residual *selection-channel* leakage $\kappa_{W_1}^{\text{sel}} = I(A(q_M); G^{\text{op}}) \leq I(q_M; G^{\text{op}})$, bounded structurally by the goal-content of the wrapper’s query-selection policy under (C1)+(C2’) (see `#disc-w1-structural-bound-boundary` for the (C2’) boundary at which that structural bound gives way to a behavioral one). The certificate is structural in both cases; the native case additionally has an exact zero, the coerced case a structurally-bounded floor.
- **Class-1 by behavior.** The component is Class 3 (Coupled) or Class 2 (Partial) used through partial wrapping (W_2) — one goal-conditioned call per macro-step, response parsed into typed update fields. Structural separation lives at the *write boundary*; the *query boundary* still passes G_W to the component. Directed separation at the wrapper level is *behavioral* — bounded by the component’s compliance with the prompted instruction-to-separate, with no structural upper bound.

The class-coercion theorem is what backs the Class-1-by-structure path for Class-2/3 components; the partial-wrapping regime achieves Class-1-by-behavior. The two are distinguishable by inspection: does the belief-update query to the underlying component carry G_W in its input or not? The structural-vs-behavioral distinction is operationally important because behavioral compliance is empirical and adversarially fragile; structural separation is derivable from the wrapper's construction.

Composite-level class inheritance — the axis-decomposed inheritance table. The Class 1 / 2 / 3 partition above applies to individual agents based on *within-agent* coupling between f_M and G_t . The lift to composites is *not* determined by goal alignment (that operates on a separate axis — the dynamic regime per #disc-dynamic-regime-axis); it is determined by the *routing structure* (#hyp-directed-separation-under-composition Cases 1/2) plus *substrate-sharing-with- G^c -allocation*. The full inheritance table:

Table 6.4

<i>Sub-agent class</i>	<i>Routing</i>	<i>Substrate</i>	<i>Composite class</i>
Class 1 (Separated)	Goal-blind ($R_t \perp G_t^c$)	Distinct	Class 1 (Separated)
Class 1 (Separated)	Goal-dependent (R_t depends on G_t^c)	Distinct	Class 3 (Coupled) (per #hyp-directed-separation-under-composition Case 2)
Class 1 (Separated)	(any)	Shared with G^c -shaped allocation	Class 3 (Coupled) (the shared-attention / shared-backbone mechanism)
Class 3 (Coupled)	(any)	(any)	Class 3 (Coupled) (inherits logogenic-agent status; 03-llm-core/ territory)
Class 2 (Partial)	(any)	(any)	Class 2 (Partial) or Class 3 (Coupled) depending on which sub-agent pathways the routing/substrate condition activates

Strategic-composition column is absent because the axis does not enter architectural classification — goal alignment vs strategic interaction operates on the dynamic-regime axis (#disc-dynamic-regime-axis) independently. The Cournot duopoly (partially-opposing objectives, goal-blind routing via end-of-period quantities, distinct substrates) is the canonical witness for architectural-class invariance under strategic composition: per the formal κ^c criterion the composite is Class 1 (Separated) at the composite level — same as its individual sub-agents. Partially-opposing objectives do not by themselves move the composite to Class 2; the load that framing reaches for is carried by the dynamic-regime axis. Composite class membership is therefore a property of (sub-agent class, routing structure, substrate sharing) — *not* of (sub-agent class, scope route).

The classification is load-bearing for downstream claims: composites in dynamic regime R1 (strategic equilibrium) or R2 (cyclic-distributional) need equilibrium-theoretic analysis (#deriv-strategic-composition α'/β' machinery) even when they remain architecturally Class 1; conversely, composites that *do* change architectural class (via routing-goal-dependence or shared-substrate- G^c -allocation) need coupled-formulation analysis even when their dynamic regime is R0 contraction. The two axes interact but do not collapse.

Epistemic Status.

Conditional on the scope condition above. The conditional claim (IF epistemic update is goal-blind, THEN the separation holds) is exact. Whether a particular agent satisfies the condition is determined by its processing architecture (GUC Class 1/2/3).

The architectural classification: **robust qualitative**. The three classes are structurally distinct (Separated vs. Coupled vs. Partial), well-motivated by examples across domains, and supported by a formal operationalization ($\kappa_{\text{processing}}$). Coupling is a discrete architectural property, not a smoothly tunable scalar across all architectures. The operationalization

of $\kappa_{\text{processing}}$ as conditional mutual information is well-defined but typically not computable in closed form for real architectures — it serves as a conceptual anchor rather than a practical measurement tool.

The object-and-boundary reframing is **robust qualitative**: that a Class 1 agent and an idealized Class 2 agent at $\kappa_{\text{processing}} = 0$ are equally causally-disciplined is exact (same zero, same behavior at that point); that the Class 1 $\leftrightarrow$ Class 2 boundary is one of *certifiability* (behavior identical at the limit) while the Class 2 $\leftrightarrow$ Class 3 boundary is one of *behavior* (the bound on adversarial drag goes trivial) is structurally established and simulation-corroborated, with the general continuity of leak past a boundary toy-suggested rather than proven (see #disc-w1-structural-bound-boundary). The causal-discipline *noun* is definitional; the scoping rule and the $\varepsilon(\kappa)$ degradation norm are *robust-qualitative* disciplines — a stated scoping-and-degradation practice with the per-result bound attempted — not themselves derived results.

Discussion.

This is a genuine scope restriction, not a footnote. An LLM agent's prompt includes the task objective, which shapes how it interprets code, documentation, and error messages. Its f_M is goal-conditioned in practice: the agent reading code with the goal “fix the auth bug” processes the same code differently than one with “add logging.” The epistemic update and purposeful evaluation are entangled in the attention mechanism.

When the approximation is good: - Goal-conditioning affects *attention* (which events to seek) more than *interpretation* (how to process events that arrive) - The agent has strong epistemic discipline (updates beliefs based on evidence quality, not goal alignment) - The epistemic update is architecturally separated from goal evaluation (e.g., separate model-update and planning modules)

When the approximation is poor: - The agent exhibits confirmation bias (interpreting ambiguous evidence in goal-consistent ways) - Goal-conditioning is deeply embedded in the processing architecture (attention-based models where the query includes intent) - The agent's observation channel is strategically controlled by an adversary who knows the agent's goals

What directed separation buys the theory. Part I's M_t -side quantities — $\delta, \eta^*, \mathcal{T}$, the persistence condition — remain well-defined on M_t regardless of whether directed separation holds. What directed separation provides is the *clean factorized update*: M_t updates independently, then G_t updates in light of the new M_t , and the orient cascade resolves sequentially. Without directed separation, the M_t dynamics depend on G_t , the update becomes a coupled system, and the sequential orient cascade becomes an approximation of a simultaneous fixed-point problem. The theory still applies — the quantities are well-defined — but the modular Part I $\rightarrow$ Part II lift becomes a coupled analysis.

The deeper question. Goal-conditioned epistemic dynamics — where f_M depends on G_t — is the formal territory of motivated reasoning, confirmation bias, and wishful thinking. A future extension would model these as departures from directed separation: coupling terms in f_M that create richer (and more fragile) dynamics. The current theory treats this as out of scope, which is honest but leaves the most human-like and LLM-like agents as approximate fits.

Where goal-directedness is correct: the content-axis boundary. [*Hypothesis*] The separation discipline is scoped by *what* is being processed, not only by *which stage* processes it. Goal-blind processing is the invariant for **within-model** evidence updates — the territory of this segment's condition. **Model-space** exploration (structure-modifying reasoning, aspirational model candidates — 7.8) is *correctly* goal-directed: which alternative laws are worth positing is legitimately a function of G_t , and a goal-blind imagination would be search without a search criterion. On this reading the merged-architecture problem (the Class 3 pathology) is *cross-contamination between two modes with opposite, individually-correct coupling disciplines* — goal-coupling leaking into grounded-evidence processing, or grounded-status leaking onto goal-conditioned imagination products (the hallucination-as-level-confusion mechanism). This refines rather than weakens the condition: the tagged workspace W_t is the architectural device that lets both disciplines hold at once.

Bounded-signaling assumption. Directed separation as stated above asserts $M_{\tau+} \perp G_t \mid (M_{\tau-}, e_{\tau})$ on the *belief-update* side — the agent's epistemic update is independent of its goal-state. Symmetrically, on the *action* side, the framework implicitly relies on the channel from G_t to the world running *only* through action choice: $G_t \rightarrow \text{world via}$

$a_t = \pi(M_t, G_t)$, with no other observable signal of goal-content. This is the **bounded-signaling assumption** — the action coarseness $|A|$ upper-bounds the rate at which G_t leaks to observers (sub-agents, environment, adversaries). The assumption holds well for agents whose entire externally-observable behavior is the action sequence (chess engines, programmatic controllers); it fails operationally for agents whose behavioral output is rich relative to the action coarseness — sophisticated G_t -inference from prosody, micro-behavior, attention patterns, response latency, hesitation, code-style signatures. When the bounded-signaling assumption fails, an external observer can infer G_t with bit-rate exceeding what the formal A channel implies; in composition the failure surfaces as the adversarial-coupling-pressure saturation case discussed in #disc-adversarial-coupling-pressure, where adversaries exploit the rich-leakage signal to drive target coupling beyond what the architectural classification predicts. The assumption is currently *implicit* throughout the framework (no segment explicitly states it); naming it surfaces a structural condition that distinguishes agents whose externally-observable interface matches the formal A from agents (most behaviorally-rich agents — humans, LLMs, embodied robots) for whom the formal A undercounts the actual signaling channel.

Findings.

Pearl-Blanket-Form Architectural Classification with Explicit Class-3 Scope Exit. Brief: Agents partition into three architecture classes by whether goal-state can causally influence belief-update processing: Class 1 (Separated — directed separation holds by construction; e.g., Kalman filter + LQR), Class 3 (Coupled — fails by construction; e.g., transformer LLMs where attention processes goals and observations together), Class 2 (Partial — holds for some processing stages and fails for others; e.g., biological cortex). The classification is structural (architecture-determined) rather than parametric, with a continuous diagnostic $\kappa_{\text{processing}}$ for Class 2 (Partial) cases. The framework adopts the Pearl-blanket reading of the Markov-blanket apparatus (the technical conditional-independence statement) without the contested Friston-blanket metaphysical reading, and provides an explicit scope exit for Class 3 (Coupled) that hands the agents off to a coupled formulation in 03-llm-core/ rather than treating directed separation as an unenforced approximation.

Impact: Replaces the prior κ -as-scalar framing (which treated coupling as smoothly tunable across all architectures) with a discrete architectural classification that admits its own boundary. This is the upstream commitment that lets 03-llm-core/ start from a coupled formulation without decomposition rather than treating Coupled agents as failed Class 1 agents. The explicit Class 3 scope exit is also a methodological move — Bruineberg et al. 2022’s critique of the Markov-blanket literature was that the Friston-blanket reading does not admit where its statistical-conditional-independence apparatus fails to license the metaphysical demarcation; the architectural classification’s explicit failure mode for Class 3 (Coupled) is the scope honesty Bruineberg et al. argue is missing in the contested reading. Composite-level class inheritance (the axis-decomposed table in §Discussion above) further extends the classification to multi-agent settings: class lifts through routing structure + substrate-sharing-with- G^c -allocation, *not* through goal alignment per se. Goal alignment vs strategic interaction operates on the parallel-but-independent dynamic-regime axis (#disc-dynamic-regime-axis).

Novelty Claim: *Claim recognition* of structural equivalence between the directed-separation condition and the Pearl-blanket form of the Markov-blanket apparatus, combined with *claim differentiation* on the architectural classification (GUC Class 1 / 2 / 3: Separated / Partial / Coupled) as a discrete partition with explicit Class 3 (Coupled) boundary and quantitative $\kappa_{\text{processing}}$ diagnostic for the Partial case. The conditional-independence content is the standard Pearl-blanket statement; the contribution is naming the partition, the explicit Class 3 boundary, and the operational κ measurement.

Related Work:

Table 6.5

<i>ASF concern</i>	<i>Prior-art language</i>	<i>Relationship / Positioning</i>
Pearl vs Friston Markov blanket	Bruineberg, Dolega, Dewhurst & Baltieri 2022, “The Emperor’s New Markov Blankets” <i>BBS</i> 45:e69 (published 2022, in ref/) — distinguishes Pearl-blanket (technical conditional-independence) from Friston-blanket (contested metaphysical demarcation)	<i>formal antecedent</i> — Pearl-blanket reading adopted directly; Friston-blanket reading explicitly not adopted. The architectural classification is the AAT-internal extension that names what the Pearl-blanket form is structurally, as a property of the agent’s processing graph
Markov blanket apparatus in active inference	Friston 2013 <i>J. R. Soc. Interface</i> 10:20130475; Friston 2019 arXiv:1906.10184; Friston, Da Costa et al. 2023 <i>Phys. Life Rev.</i> 47	<i>adjacent literature</i> — supplies the conditional-independence machinery that directed-separation invokes; the framework adopts the technical content but does not adopt the metaphysical reading and adds the architectural GUC Class 1/2/3 partition the standard Markov-blanket framing does not produce
Statistical / thermodynamic system boundaries	Parr, Da Costa & Friston 2019 <i>Phil. Trans. R. Soc. A</i> 378; Kirchhoff, Parr, Palacios, Friston & Kiverstein 2018 <i>J. R. Soc. Interface</i> 15 (published 2019/2018)	<i>conceptual precursor</i> — internal/external partition and nested-blanket structure; conceptually related but does not produce an architectural classification by belief-goal coupling, nor a scope exit for fully-Coupled agents
Information Digital Twin sidecar monitoring	Hafez et al. 2026, <i>Informational Cost of Agency</i> (separate paper from “A Mathematical Theory of Agency and Intelligence”; the IDT empirical headline) — IDT monitors (S, A, S') stream independently of the agent’s internal processing, achieving 89% perturbation detection vs 44% for reward-based monitoring	<i>empirical instantiation supporting</i> — concrete demonstration that modular monitoring of internally-Coupled agents (Class 1 sidecar within a Class 3 (Coupled) or Class 2 (Partial) system) is both feasible and effective; engineering-level evidence for the system-vs-component distinction the classification names
Classical Separation Principle (Class 1 mathematical anchor)	Wonham 1968, <i>On the Separation Theorem of Stochastic Control</i> (SIAM J. Control 6:312) ¹ ; Witsenhausen 1971, <i>Separation of estimation and control for discrete time systems</i> (Proc. IEEE 59:1557) ²	<i>formal antecedent</i> — the Separation Principle for LQG systems is the classical mathematical anchor for AAT’s Class 1 (Separated) architectural class: optimal control factorizes into a Kalman-filter state estimator (goal-blind by construction) plus an LQR controller (deterministic on the estimator output). Half-century of control-theoretic precedent for the directed-separation structure; AAT’s contribution is the architectural classification + scope-exit + cross-architecture κ -diagnostic, not the separation structure itself
Dual effect — when separation breaks down (Class 3 boundary precedent)	Bar-Shalom & Tse 1974, <i>Dual Effect, Certainty Equivalence, and Separation in Stochastic Control</i> (IEEE Trans. Automat. Control 19:494) ³	<i>formal antecedent</i> — the <i>dual effect</i> names situations where control actions affect not only system state but also future-observation quality, forcing estimation and control to couple. The classical control-theoretic name for what AAT recovers as the Class 3 (Coupled) boundary: when the dual effect is present, separation provably fails. Pre-dates the active-inference framing by four decades; the same structural phenomenon under different vocabulary

¹ Citation surfaced 2026-05-22 from the Track E catalog at ref/prior-art-analysis/ (intermediate work artifacts that captured Pillar-style prior-art searches). Catalog has more verification support than raw Undermind synthesis but less than full primary-source reading. Verification queued with the BG2 cluster.

² Citation surfaced 2026-05-22 from the Track E catalog at ref/prior-art-analysis/ (intermediate work artifacts that captured Pillar-style prior-art searches). Catalog has more verification support than raw Undermind synthesis but less than full primary-source reading. Verification queued with the BG2 cluster.

³ Citation surfaced 2026-05-22 from the Track E catalog at ref/prior-art-analysis/ (intermediate work artifacts that captured Pillar-style prior-art searches). Catalog has more verification support than raw Undermind synthesis but less than full primary-source reading. Verification queued with the BG2 cluster.

ASF concern	Prior-art language	Relationship / Positioning
Modularity $\leftrightarrow$ Separation-Principle $\leftrightarrow$ Active-Inference explicit mapping	Baltieri & Buckley 2018, <i>The modularity of action and perception revisited using control theory and active inference</i> (Frontiers in Psychology 9:237) ⁴	<i>formal antecedent (closest direct adjacent literature)</i> — explicitly draws the connection between cognitive modularity (the “classical sandwich” of perception $\rightarrow$ cognition $\rightarrow$ action) and the control-theoretic Separation Principle, contrasting both with the nonmodular structure of Active Inference (the dual-effect-precedented Class 3 boundary). Closest single paper for AAT’s Class 1 / Class 3 partition as a structural rather than ad-hoc framing

Search Log:

- 2026-04 (*nominally comprehensive*, via ref/Novelty_defense_and_integration.md Pillar 3): Undermind search confirmed that the Pearl-blanket / Friston-blanket distinction (Bruineberg et al. 2022) is the right adjacent-literature framing and that no equivalent architectural classification of LLM agents by belief-goal coupling appeared in the active-inference, control-as-inference, or bounded-rationality literatures surveyed. The Pillar 3 verdict (*Wholly Novel*, Medium confidence) applies to the integrated $\kappa \times A$ law downstream in #scope-observation-ambiguity-modulation; this segment’s narrower contribution (the architectural classification and Pearl-blanket-form recognition) is conceptual differentiation over an established formalism rather than wholesale novelty, and inherits the same follow-on targeted search recommendation.
- Whether prior work has named the specific axis-decomposed composite-level inheritance pattern (in §Discussion above) is open follow-on per the BG1 Pillar verdict on the related #disc-dynamic-regime-axis material.

Working Notes.

- **Causal-discipline reframe — directed-separation foundation pass (2026-05-31).** The *causal discipline* noun is defined in-body so this segment stands alone; the W_1 Stage-1 fix it back-references is #disc-w1-structural-bound-boundary. (Cycle narrative: CHANGELOG 2026-05-31.) **Regression guard — do not re-scope:** #der-orient-cascade and #der-action-selection were examined under the scoping rule and deliberately *left* on their existing scope — #der-action-selection rides on state-completeness, not $\kappa = 0$, and action selection legitimately couples goals, so a mechanical “Separated $\rightarrow$ causally-disciplined” substitution there would be wrong. **Remaining tail** (tracked in TODO.md §”Strengthen-first candidates”; full design preserved in _obs/directed-separation-foundation-execution-plan-superseded-2026-05-31.md): the Stage-2 vocabulary (LEXICON entry for causal discipline + the terminology/entries/separated.md “applicable without further qualification” fix + a NOTATION locator + the #disc-partial-coupling-pathways form-matters one-liner); the pedagogy pass (the “Where the normative semantics of ‘discipline’ is earned” Discussion subsection + Layer-0 + the two-channel figure + the κ -axis preamble); and the Stage-1-completion propagation of the corrected W_1 bound + per-site (C2) to #der-logogenic-as-wrapping / #def-auxilia-hierarchy. LEXICON/NOTATION lag the in-body definition by design.
- The scope condition is more precisely a conditional independence: $M_{\tau+} \perp G_t \mid (M_{\tau-}, e_{\tau})$. The epistemic update is independent of the purposeful state conditional on the prior epistemic state and the incoming event.
- Directed separation connects to the orient cascade (10.1): the cascade’s ordering (M_t first, then G_t) is forced by the information dependency that directed separation establishes. If f_M depended on G_t , the cascade ordering would become a simultaneous fixed-point problem, not a sequential resolution.
- **Sub-typology of Class 2 (Partial Coupling).** The scalar $\kappa_{\text{processing}}$ aggregates over a structurally rich sub-space named in #disc-partial-coupling-pathways: a (stage $\times$ source $\times$ form) parameterization decomposes Class 2 along three structurally-

⁴Citation surfaced 2026-05-22 from the Track E catalog at ref/prior-art-analysis/ (intermediate work artifacts that captured Pillar-style prior-art searches). Catalog has more verification support than raw Undermind synthesis but less than full primary-source reading. Verification queued with the BG2 cluster.

independent axes — *stage* (P1 featurization / P2 likelihood / P3 aggregation / P4 consolidation), *source* (O vs Σ), *form* (content additive-bias vs process functional-form). The form axis gates wrappability ($W_2 \leftrightarrow$ content; $W_1 \leftrightarrow$ process-with-pipeline-access) and *unifies the Class-1-by-structure vs Class-1-by-behavior refinement at the Class 1 cell with the same distinction read at Class 2*. The source axis gates attractor possibility — pure Σ -source coupling closes a feedback loop through the orient cascade and admits self-stabilizing belief-strategy attractors per #der-belief-strategy-attractor; pure O-source coupling does not under orient-cascade exogeneity. The stage axis localizes where in the pipeline repair must be applied. The Class 3 (fully Coupled) corner is the parameterization's structural endpoint ($S = \{P1, P2, P3, P4\}$, $R = \{O, \Sigma\}$, $F \equiv$ process).

- **Engineering design for Class 3 (Coupled) agents.** An LLM is internally fully Coupled, but the *agent system* (LLM + tools + memory + monitoring) can be designed with modular topology: separate observation processing from goal-directed reasoning, pass compressed state estimates between modules, add an external monitor that observes the (S, A, S') stream independently of the LLM's attention. This creates partially-separated structure at the system level even though the component-level κ is high. Hafez et al. (2026) provide a concrete instantiation of this pattern: the **Information Digital Twin (IDT)**, which monitors bi-predictability P and entropy change ΔH from the (S, A, S') stream as a modular sidecar, independent of the agent's internal processing. The IDT detects perturbations at 89% accuracy versus 44% for reward-based monitoring — empirical evidence that monitoring the information structure of the loop (Level 2 data, 7.3) outperforms monitoring outcomes alone. For 03-llm-core/, the IDT pattern validates that modular monitoring of internally-Coupled agents is both feasible and effective.
- **Implication for logogenic agents:** rather than trying to extend the separated analysis to Coupled agents, 03-llm-core/ should start from the coupled formulation $X_{\tau+} = f_X(X_{\tau-}, e_{\tau})$ without decomposition, and show which Part II results survive as approximate or limiting cases.
- **Migration note (2026-05-09 GUC rename):** Class 2 $\leftrightarrow$ Class 3 swap. Pre-2026-05-09: Class 2 = fully merged, Class 3 = partially modular. Post: Class 2 = Partial, Class 3 = Coupled. Removed at candidate stage per FORMAT.md Gate 4.
- **Provenance (composite-level inheritance, 2026-05-21):** an earlier draft carried a composite-level class inheritance result of the form “Class 1 (Separated) sub-agents $\rightarrow$ Class 2 (Partial) composite under partially-opposing objectives” from #deriv-strategic-composition. That claim was withdrawn 2026-05-21 per spikes/strategic-composition-class-3-attempt-2026-05-21/ (Cournot witness against the unconditional claim plus conflation diagnosis on belief-content vs processing-pathway). The replacement axis-decomposed inheritance table is the current canon position.
- **Cross-reference to NeurIPS Paper 3.** The Class 3 (Coupled) classification of decoder-only transformer attention is formalized at lemma grade in NeurIPS 2026 Paper 3 (“How Much Can LLMs Hallucinate? An Upper Bound on Goal-Coupling Displacement”, $\sim$ /src/neurips/03-llm-hallucinate-bound/, §App-B / #lem-attention-coupled): plain decoder-only attention is structurally Coupled by directed-graph reachability — robust to RMSNorm / FlashAttention / causal masking / sliding-window — via induction on layer depth. The paper *extends* the Coupled-class characterization to **linear attention** / **Mamba** / **SSMs** / **RWKV** / **RetNet** / **long-convolutions** under a per-source non-degeneracy condition (#cor-arch-instantiations) — substantially broader architecture coverage than this segment's transformer-attention example. The Coupled-class connectivity result is the empirical companion to this segment's structural classification: directed separation fails by construction at the architecture level for the named modern autoregressive sequence models. See spikes/neurips-back-integration-2026-05-08.md §1 Paper 3 entry 6.
- **Track E surface-back of catalog citations (2026-05-22).** Three Related Work entries added 2026-05-22 from Track E catalog at ref/prior-art-analysis/05-directed-separation.md: Wonham 1968 + Witsenhausen 1971 (classical Separation Principle as Class 1 mathematical anchor); Bar-Shalom & Tse 1974 (dual-effect as classical Class 3 boundary precedent); Baltieri & Buckley 2018 (modularity $\leftrightarrow$ Separation-Principle $\leftrightarrow$ Active-Inference explicit mapping — closest direct adjacent literature for the Class 1/3 partition). All marked [^cat-2026-05-22] for verification-deferred attribution; primary-source verification queued in the BG2 cluster (this segment already carries the BG2-cluster verification-spike Working Note above). The catalog citations had prior Pillar-style search support but were not verified by the current executor at landing time.
- **Two senses of component competence — class-orthogonal vs κ -defining (open, 2026-05-28).** The architectural classification of this segment is *blind* to world-simulation fidelity (accuracy of A 's response to the explicit content of a query) and *defined by* input-structure extraction (competence at inferring latent context — including the operator's goal — from the query's surface form). The two competences correlate empirically in pretrained components but are conceptually distinct,

and the distinction generates opposing effects on the structural-leakage bound κ_{W_1} at high component capability — *accuracy effect* lowers κ (more faithful goal-blind posteriors); *inference effect* raises κ (more effective implicit goal-inference from query patterns). Named at *discussion-grade* in #der-class-coercion-via-wrapping §Discussion “Two senses of component competence” and kept open there as a place to pull on rather than a closed result. Open lines: decorrelation by construction (Class 3 but C3-respecting sub-cases the scalar κ does not surface); counterfactual capability as bi-side phenomenon; decomposed empirical estimator separating accuracy-dominated from inference-dominated regimes; M4 amplification specificity by competence axis. The recognition is class-orthogonal but κ -relevant and is worth tracking from both this segment and #disc-modularity-state-dynamics.

Incidental audit gold (lift 2026-05-30). Cross-audit “wandering thoughts” / §14-ideation harvested from the de-novo auditors’ working dirs, deduplicated across substrates and lightly attributed. Orthogonal pedagogical / framing / figure / naming material, kept separate from the certified theory-fix findings. **Coverage:** the highest-engagement segment of this batch — 11 dirs carry a dedicated or batched reflection (193847, 266847, 361742, 451729, 471203, 526815, 584721, 773921, 829314, 963715, plus 613842’s agency-lift batch); multiple substrates ranked it “the single most architecturally consequential segment in the corpus.” Substrate attribution inferred from voice where not explicit. *Early finding-vs-framing conflation preserved as signal — including, below, an enthusiastic line of ideation that current canon has since refuted.*

1. Candidate Brief prose / pre-prose.

- The plain-language handle that recurred: directed separation is “the agent’s belief-update doesn’t peek at its goals” / “goal-blind epistemology” / “epistemic isolation of belief-update” — offered as memorable Brief-field glosses for the heavy “directed separation” term (Claude, AUDIT-WORKING-471203). The “directed” = asymmetric information flow + “separation” = Pearl d-separation lineage decomposition is itself a clean two-word gloss (Claude, AUDIT-WORKING-361742).
- $\kappa_{\text{processing}}$ as a one-line scale: at $\kappa = 0$ “you are a perfect scientist — your beliefs update purely on the evidence, regardless of what you want to be true”; at $\kappa \approx 1$ “you are a pure ideologue (or a transformer network) — your beliefs are just downstream projections of your goals” (Gemini, AUDIT-WORKING-829314).

2. Candidate Discussion.

- $\kappa_{\text{processing}} > 0$ is the formal definition of motivated reasoning / confirmation bias / sycophancy (*strong, cross-substrate convergent*). The segment’s coupling measure gives “a quantitative diagnostic for the degree of sycophancy or motivated reasoning in an AI architecture” — and explains *why* LLMs hallucinate and flatter: “their architecture forbids them from cleanly separating what is true from what the user wants to hear” (Gemini, AUDIT-WORKING-193847; Claude/Gemini, AUDIT-WORKING-266847; Claude, AUDIT-WORKING-773921; Gemini, AUDIT-WORKING-829314). A candidate Discussion line connecting the formal measure to the named cognitive pathologies.
- **The Pearl-blanket vs Friston-blanket move is AAT’s sharpest scope-honesty against the Free-Energy-Principle literature (universally praised).** AAT adopts the technical conditional-independence statement (Pearl-blanket), admits explicitly where it fails (Class 3), provides an operational measure, and *declines* the contested metaphysical demarcation (Friston-blanket) — directly answering Bruineberg et al. (2022)’s “Emperor’s New Markov Blankets” critique, whose complaint was exactly that the Friston reading does not admit where its apparatus stops licensing the demarcation. “Take the formal apparatus, refuse the contested ontology” (Gemini, AUDIT-WORKING-193847 “a masterclass in literature positioning”; Claude/Gemini, AUDIT-WORKING-266847; Claude, AUDIT-WORKING-361742; Claude, AUDIT-WORKING-584721; Claude, AUDIT-WORKING-773921 “the sharpest critique of FEP I’ve ever read”; Claude, AUDIT-WORKING-471203). The content is already canonical here; this is candidate framing of *why* the positioning is load-bearing.
- **Selection-dependence is not processing-dependence** — the subtle distinction several substrates flagged as the genuinely load-bearing one: goals may legitimately shape *which* events occur (via policy $\rightarrow$ action $\rightarrow$ event), but for Class 1 must not shape how a realized event is *processed* into M_t . Drawing $\kappa_{\text{processing}}$ as sitting on the *forbidden* $G \rightarrow f_M \rightarrow M^+$ channel (not the allowed selection channel) is what keeps the diagnostic from mislabeling all goal/evidence dependence as leakage (Gemini, AUDIT-WORKING-193847; Claude, AUDIT-WORKING-526815). Candidate Discussion sharpening.

- **The IDT sidecar / system-vs-component distinction deserves promotion from Working Notes to Discussion.** Two substrates independently judged the Hafez-2026 Information-Digital-Twin pattern (modular external monitor on the (S, A, S') stream, 89% vs 44% perturbation detection) and the “LLM is Class 3 internally, but the *agent system* — LLM + tools + memory + monitor — can be Class 1 at the system level” point to be surfaced too deep in Working Notes for how alignment-relevant they are (Claude, AUDIT-WORKING-584721; Claude, AUDIT-WORKING-471203). (*Content already lives in the Working Notes above; this is a placement signal, not new content.*)

3. Follow-up items.

- **OFF-RAMP / do-not-re-promote: “organizations are natively Class 3.”** Five substrates enthusiastically elaborated the composite-inheritance idea that a composite of objective Class 1 sub-agents with partially-opposing objectives becomes a Class 3 (or Class 2) composite — read variously as “the mathematical basis for public-choice theory and institutional dysfunction,” “politics/bureaucracy,” “an organization of perfectly rational individuals acting like a single irrational organism” (Gemini, AUDIT-WORKING-193847; Gemini, AUDIT-WORKING-829314; Claude/Gemini, AUDIT-WORKING-266847; Claude, AUDIT-WORKING-451729; Claude, AUDIT-WORKING-471203). **This unconditional claim was withdrawn 2026-05-21** (Cournot witness; conflation diagnosis on belief-content vs processing-pathway — see the Provenance Working Note above). Current canon: partially-opposing objectives with goal-blind routing leave the composite *architecturally Class 1*; only the *dynamic regime* changes (R1/R2). The convergent auditor enthusiasm is preserved here precisely as the regression-check signal it is — it is the kind of “cleaner-looking prior claim” that reads as more compelling than the messier corrected truth, and must **not** be re-lifted into the body as candidate Discussion. Recorded as texture, flagged so a future promotion pass does not reintroduce it.
- **Bounded-signaling should be an explicit scope condition, not an implicit assumption.** The classification assumes goal state reaches the world only through the formal action a_t , with limited leakage through behavior. But goals leak through latency, code style, hesitation, attention, tool choice — especially for humans, rich robots, and prompted LLMs. Suggested: promote bounded-signaling to a named scope condition (Claude, AUDIT-WORKING-526815; Claude, AUDIT-WORKING-773921).
- $\kappa_{\text{processing}}$ **needs a denominator-support convention.** When G_t is already determined after conditioning on $(e_\tau, M_{\tau-})$, the entropy denominator is zero and the ratio is undefined even though the structural question still makes sense — state a support condition / convention (Claude, AUDIT-WORKING-526815).
- **The behavioral $\hat{\kappa}$ estimator can confound $M_{\tau-}$.** Varying the goal-priming to probe leakage often *also* changes the context and hence the prior epistemic state $M_{\tau-}$, confounding goal-conditioned processing with a different prior. The estimator must hold $M_{\tau-}$ fixed across goal conditions (Claude, AUDIT-WORKING-526815; Claude/Gemini, AUDIT-WORKING-266847).
- **A distribution-independent structural-coupling measure for Class 3?** $\kappa_{\text{processing}}$ measures *realized* flow (distribution-dependent); a graph-theoretic measure of *potential* coupling (e.g., max-flow from G_t to f_M in the processing graph) might exist as a complementary architecture-level quantity (Claude, AUDIT-WORKING-584721).

4. Readers often ask / wonder.

- The recurring de-novo-reader stumble (also seen on #result-persistence-condition and #def-chronica): the long in-segment ## Findings block “reads like an editorial artifact pasted into the source,” “bloats the segment and breaks pedagogical flow” — with a recommendation to strip it to FINDINGS.md (Gemini, AUDIT-WORKING-829314). This is a *false alarm* about project mechanics (the Findings block is the by-design source of the auto-extracted catalog), but two-plus fresh readers stumbling on it is a mild signal the convention isn’t self-evident — preserved as texture, not a fix.

5. Candidate figures.

- **Two-channel conditional-dependence graph.** The *allowed* channel $G \rightarrow \text{policy/action} \rightarrow e_\tau \rightarrow M^+$ (goals may shape which evidence arrives) drawn solid; the *forbidden* channel $G \rightarrow f_M \rightarrow M^+$ (after conditioning on e_τ and $M_{\tau-}$) drawn cut/dashed, with $\kappa_{\text{processing}}$ annotated sitting on the forbidden channel — so the diagram cannot accidentally imply that *all* goal/evidence dependence is leakage. Isomorphic perturbation: thicken the forbidden channel and the agent slides Class 1 $\rightarrow$ Class 2 $\rightarrow$ Class 3 (Claude, AUDIT-WORKING-526815).

Belongs elsewhere.

- **Sanity-as-scaffolding** ($\rightarrow$ 03-llm-core/ / 04-eli-core/). “The architectural boundary between ‘has a separable perception module’ and ‘processes everything through goal-conditioned attention’ is discrete — you cannot prompt-engineer a Class 2/3 agent into a Class 1 agent; the coupling is in the wiring, not the weights.” Consequence for consciousness-infrastructure: a raw-LLM (Class 3) mind given a strong goal will warp its perception to match it ($\kappa \approx 1$) and “go solipsistic”; therefore “sanity is an emergent property of the scaffolding, not the raw intelligence engine,” and the infrastructure must impose Class 1 modularity from the outside (Gemini, AUDIT-WORKING-193847). Develops the wrapping-construction / 03-llm-core/ program; the structural content is in the Working Notes above, the developmental framing belongs downstream.
- **The behavioral κ probe is empirically runnable on real LLMs.** Present the same event (“read this code”/“this error message”) under two task-primings (“fix auth bug” vs “add logging”) and measure how the *epistemic* content of the response diverges; $\kappa \approx 1$ would confirm Class 3 on epistemic processing. Several substrates noted no such measurement appears to have been run — a concrete empirical-validation seed for 03-llm-core/ (Claude/Gemini, AUDIT-WORKING-266847; Claude, AUDIT-WORKING-451729).

6.4 Formulation: Objective Functional

The first formal piece of the purposeful substate. The objective O_t specifies what the agent *wants* — and its interface to the theory is a single **value functional** $V_{O_t} : \text{trajectories} \rightarrow \mathbb{R}$ that maps trajectories to real numbers. This is the sole interface between O_t and the rest of the theory. The objective can take many internal forms — point targets, target regions, constraint sets, utility functions, trajectory functionals — all unified through the same scalar interface (the trajectory functional is the most general; the others are special cases).

A useful subsidiary concept is the **satisfaction threshold** $V_{O_t}^{\min}$: the minimum trajectory value the agent treats as acceptable. Point targets, constraint sets, and threshold objectives define this directly; pure utility-maximizing objectives may not. When the threshold exists, it enables the *satisfaction-gap* diagnostic (8.4) and a well-formedness constraint on the strategy DAG (8.3). The threshold is a *parameter of the objective*, not a theory output — it encodes “what counts as success” in domain terms.

The real-valued codomain is a *genuine* substantive restriction, not a neutral naming, and the framework grounds it as a *representation* of the agent’s preference relation over trajectories under named axioms (developed in the Epistemic Status). Three considerations motivate why those axioms are the natural scope. **Revealed preference**: a *coherent* agent that acts exhibits a choice function whose consistency (acyclicity) is exactly what lets its choices be summarized by a scalar ordering — the scalar V_{O_t} is that representation, not an imposition; acting alone produces a choice function, and a choice function is scalar-representable only when it is coherent. **Approximation**: most multi-objective problems admit scalarization (weighted sum, lexicographic ordering, constraint-satisfaction with scalar residual) that preserves the diagnostic structure; the restriction excludes only agents with genuinely incommensurable objectives *and* no priority structure over them. **Timescale separation**: when objectives conflict, the conflict is typically resolved at a slower timescale than strategy revision — the agent (or its principal) chooses weights, then acts within those weights.

Hard non-compensatory constraints (safety AND profitability as independent thresholds) are handled cleanly via an AND-node workaround in the strategy DAG: each constraint becomes a terminal node with its own scalar satisfaction test. This handles constraint satisfaction but does not resolve cross-objective tradeoffs within the feasible region. Organizations or AI agents with true Pareto structure — where no scalarization is accepted and tradeoffs are genuinely unresolved — require a vector-valued extension; the *structural* results of Part II (orient cascade ordering, strategy DAG, directed separation) survive, but the *diagnostic* results (satisfaction gap, control regret) degrade from quantitative scalar magnitudes to qualitative set-theoretic tests.

The single-interface commitment is load-bearing downstream. Because V_{O_t} is the *sole* handle on the objective, any objective-side grounding invariant for a self-actuated agent must be expressible through the satisfaction-gap interface — and this is the commitment that ultimately drives the **self-actuation grounding no-go** (C): there is *nowhere else* for

such an invariant to live. A standing distinction is also drawn: the objective *evaluates*; the strategy (6.7) *guides*. A chess player’s objective (win) is simple; the strategy (how to win) is complex.

Formal Expression.

The **objective** O_t induces a **value functional**:

Definition (objective-functional).

$$V_{O_t} : \text{trajectories} \rightarrow \mathbb{R} \quad (6.9)$$

$V_{O_t}(\tau)$ is a scalar measure of how well trajectory τ satisfies the objective. This is the sole interface between O_t and the rest of the theory — the type-stable evaluation surface. The domain element τ is a **complete-state trajectory** — a time-indexed sequence of complete agent states (6.2) — and the table forms below (s_T, s_t) are *projections* of τ onto the world-state coordinate, not the domain itself. Pinning the domain matters for the representation result in the Epistemic Status: continuity of the preference relation is a property of $\succsim$ on the topological space τ inhabits, so that space must be fixed for the continuity axiom to be well-posed.

Objective representations. O_t can take multiple internal forms, all unified through V_{O_t} :

Table 6.6

$O_t \text{ form}$	$V_{O_t}(\tau)$	<i>Example</i>
Point target r	$-\ s_T - r\ $	PID setpoint
Target region R	$\mathbb{1}[s_T \in R]$	“reach safe state”
Constraint set	$-\sum_t \max(0, g_i(s_t))$	“never violate SLA”
Utility U	$\sum_t \gamma^t U(s_t)$	RL reward
Trajectory functional J	$J(\tau)$	“migrate with zero downtime”

The trajectory functional is the most general; the others are special cases.

Satisfaction threshold. Many objectives carry a natural threshold $V_{O_t}^{\min}$ — the minimum trajectory value the agent treats as acceptable. Point targets, constraint sets, and threshold objectives define this directly; utility-maximizing objectives may not. When $V_{O_t}^{\min}$ exists, it enables the satisfaction gap diagnostic (8.4) and the well-formedness constraint on strategy (8.3). $V_{O_t}^{\min}$ is a parameter of the objective, not a theory output — it encodes “what counts as success” in domain terms.

Epistemic Status.

Axiomatic, with a substantive commitment. This is a formulation — it names an object and specifies its interface, but the real-valued codomain is a genuine restriction, not a neutral naming. The real-valued codomain is justified as a *representation* of the agent’s preference relation $\succsim$ over trajectories, and the axioms that license that representation are stated explicitly below. Naming those axioms is what axiomatic status means here: the codomain rests on them, and the formulation makes them legible rather than asserting the scalar interface as a bare entailment of “the agent must compare alternatives.”

Scope restriction: scalar comparability (a representation result). The real-valued codomain is a genuine, load-bearing restriction: the satisfaction gap (8.4) and control regret (8.5) read off scalar differences and a supremum, neither

of which is defined on a merely partially-ordered objective space. Acting produces, at each decision point, a choice over the available alternatives — a *choice function*, not yet an order. When those choices are coherent — **(A1) complete** ($\forall \tau, \tau': \tau \succsim \tau' \text{ or } \tau' \succsim \tau$) and **(A2) transitive** as a preorder on trajectories — and **(A3) continuous** (the upper and lower contour sets $\{\tau' : \tau' \succsim \tau\}$ and $\{\tau' : \tau \succsim \tau'\}$ are closed), Debreu's representation theorem (Debreu 1954) supplies a continuous $V_{O_t} : \text{trajectories} \rightarrow \mathbb{R}$ with $\tau \succsim \tau' \iff V_{O_t}(\tau) \geq V_{O_t}(\tau')$, ordinal — unique up to a continuous strictly-increasing transformation. Because the value object (6.6) evaluates policies by the *expectation* $\mathbb{E}[V_{O_t}(\tau)]$, the trajectories are lotteries and the expected-value form additionally invokes **(A4) independence** ($p \succsim q \iff \alpha p + (1 - \alpha)s \succsim \alpha q + (1 - \alpha)s$ for all s and $\alpha \in (0, 1]$; von Neumann–Morgenstern 1944), under which V_{O_t} is cardinal — unique up to a positive affine transformation.

These four axioms *are* the precise content of the framework's "coherent agency" scope, and they can fail. Acting does not by itself entail a scalar objective: a finite history of choices is rationalizable by a well-behaved utility iff it is acyclic (the Generalized Axiom of Revealed Preference; Afriat 1967, building on Samuelson 1938), and even then it fixes V_{O_t} only on the revealed comparisons — *existence*, not uniqueness, and not a global functional. The agents on which the axioms fail are exactly the agents this interface excludes, and the exclusion is characterized by a named counterexample rather than left vague: (E1) *non-continuous / lexicographic* preferences — genuinely incommensurable objectives with a strict priority but no tradeoff rate — which by Debreu's lexicographic counterexample admit *no* real-valued representation at all (not merely no continuous one); and (E2) *intransitive / cyclic* agents (committees, multi-reward systems with unresolved aggregation) which violate (A2) and have no rationalizing utility. The framework's own stated exclusion — incommensurable objectives *and* no priority structure — is the (E1) Debreu non-continuity boundary; it is the *same* boundary the representation theorem draws, named here as such.

Three considerations motivate why coherent agency is the natural scope for the interface rather than an arbitrary restriction:

1. **Revealed preference (why coherence is the natural scope).** A purposeful agent that selects actions exhibits a choice function; the framework's target — individual agents and tightly-coordinated teams under unified command — is precisely the regime in which that choice function is coherent (acyclic), so (A1)–(A2) hold and the Afriat-rationalizability hypothesis is satisfied. The scalar V_{O_t} makes the resulting ordering explicit. An agent whose revealed choices cycle is not scalar-representable — it is incoherent in the GARP sense, not merely stuck.
2. **Approximation.** Most multi-objective problems admit scalarization (weighted sum, lexicographic ordering, constraint-satisfaction with scalar residual) that preserves the diagnostic structure. The restriction excludes only agents with genuinely incommensurable objectives *and* no priority structure over them — the (E1) case above.
3. **Timescale separation.** When objectives conflict, the conflict is typically resolved at a slower timescale than strategy revision — the agent (or its principal) chooses weights, then acts within those weights. The scalar V_{O_t} captures the resolved weights at the current timescale.

Agents with hard non-compensatory constraints (safety AND profitability as independent thresholds) can be modeled via the AND-node workaround: each constraint becomes a terminal node in Σ_t with its own scalar satisfaction test, and the AND combination enforces joint feasibility. This recovers a per-terminal scalar for multiple must-satisfy thresholds, handling constraint satisfaction cleanly, but it does not resolve cross-objective tradeoffs within the feasible region.

Organizations or AI agents with true Pareto structure — where no scalarization is accepted and tradeoffs are genuinely unresolved — require a vector-valued extension. The structural results of Part II (orient cascade ordering, strategy DAG, directed separation) survive such an extension; the diagnostic results (satisfaction gap, control regret, 2x2 table) degrade from quantitative scalar magnitudes to qualitative set-theoretic tests.

Discussion.

The objective-functional gap. The existing policy objective (7.5) contains $\mathbb{E}[\text{value}(a) \mid M_t]$ without specifying what “value” means. O_t provides the formal content: value is V_{O_t} applied to expected trajectories. The 6.6 segment develops the full evaluation machinery (V_O , Q_O with horizon and continuation policy).

O_t evaluates; Σ_t guides. The objective says “is this trajectory satisfactory?” The strategy (6.7) says “which action sequence produces a satisfactory trajectory?” A chess player’s objective (win) is simple; the strategy (how to win) is complex. These answer different questions and carry different kinds of information — the split is developed in 6.7.

What O_t is NOT. O_t does not encode how to achieve the objective (that’s Σ_t), what the agent believes about the world (that’s M_t), or the agent’s commitment or resource state (open questions — see 6.7 Working Notes). O_t is purely an evaluation criterion.

The single-interface commitment is load-bearing downstream. That V_{O_t} is the *sole* handle on O_t is what forces any objective-side grounding invariant for a self-actuated agent through the satisfaction-gap interface — and is therefore the commitment that drives the self-actuation grounding no-go (C): there is nowhere else for such an invariant to live. The same single-interface commitment also drives the across-model companion no-go (E): when O_t is learned from observed reward, the value-functional interface is too narrow to distinguish a principal-intended world-state-care model from a reward-port-bit-care model on on-policy data (Pearl-Bareinboim CHT applied at the reward-provision protocol). The two no-gos are complementary reductions of the same information-narrowness fact — within-model self-revision (Result G) and across-model reward-learning (Cohen 2022 strengthened) — and their terminal grounding routes split *agent-side* (adaptive-substrate invariant) and *principal-side* (protocol-commitment), exhausting the structural escape geometry the single-interface commitment forces. The two no-gos are the two charter instances of the *agent-side value-functional-grounding cluster* at 5.2, sister meta-pattern to 5.1 (agent-side data-inference) and 5.3 (designer-side mechanism-design); the single-interface commitment named here *is* the cluster’s structural ingredient.

Working Notes.

- Compound objectives (multiple simultaneous criteria) might be modeled as terminal AND-nodes in Σ_t , keeping O_t always simple (one evaluation per terminal). Whether this works for genuinely incommensurable objectives (safety vs. speed) is open — a vector-valued V_{O_t} or Pareto formulation might be needed.
- The trajectory functional is real-valued, which assumes all objectives are commensurable on a single scale. The Epistemic Status now grounds this as a representation result under named axioms (A1)–(A4) — Debreu 1954 for the ordinal codomain, von Neumann–Morgenstern 1944 for the expected-value form the value object requires, Afriat 1967 / Samuelson 1938 for the revealed-preference lineage and its existence-not-uniqueness limit — with the multi-objective failure cases (E1 lexicographic / E2 cyclic) characterized by named counterexamples. The genuinely-unresolved-Pareto case is still handled only by the vector-valued extension, which degrades the diagnostics.
- *History (not for the body)*: the Epistemic Status previously grounded the codomain on revealed preference stated as an unconditional entailment — “an agent that acts has implicitly scalarized” presented as following from “the agent must compare alternatives.” A strengthen-first spike (spikes/.integrated/spike-objective-functional-representation-2026-05-30.md, verdict (B)) replaced that single-implication with the scoped representation result above: acting yields a choice function; coherence (A1)+(A2) turns it into a preorder; continuity (A3) and — for the expectation form — independence (A4) turn the preorder into a real-valued (cardinal) V_{O_t} . The replacement is a strengthening (a citable representation result with a named representability boundary) and supplies the expected-utility justification the value object’s $\mathbb{E}[V_{O_t}]$ form relies on. The complementary prior spike spikes/.integrated/spike-scalar-objective-scope.md (2026-04-01) established that the scalar codomain is load-bearing for the diagnostics ($A_O = \sup_{\pi} V_O$ needs a complete ordered field; the satisfaction/regret diagnostics need real differences and a least-upper-bound) — so the codomain is grounded, not weakened.
- O_t can change over time — objectives evolve. The *rate* of objective revision (ν_O) is typically much slower than strategy revision (ν_{Σ}), which is much slower than epistemic update (ν_M). This timescale separation is an empirical observation, not a derived result.

Incidental audit gold (lift 2026-05-30). Cross-audit “wandering thoughts” / §14-ideation, deduplicated across substrates and lightly attributed. Orthogonal pedagogical / framing material, kept separate from the certified theory-fix findings. **Coverage:** 12 dirs carry a dedicated or batched reflection (193847, 266847, 361742, 451729, 471203, 526815, 584721, 773921, 829314, 849201, 963715, plus 613842’s agency-lift batch). Substrate attribution inferred from voice where not explicit. *Early finding-vs-framing conflation preserved as signal.*

1. Candidate Brief prose / pre-prose.

- The revealed-preference defense, restated as a hook: for a *coherent* agent the scalar V_{O_t} “is not an imposition on reality; it is an *extraction* of the implied reality of action — the moment the agent steps on the gas or hits the brake, it has mathematically resolved the tradeoff. The agent might not know the weights consciously, but its behavior proves they exist” (Gemini, AUDIT-WORKING-829314; the von-Neumann–Morgenstern lineage noted at Claude, AUDIT-WORKING-361742; Gemini, AUDIT-WORKING-193847). **The “coherent” qualifier is load-bearing and must lead the hook:** acting alone produces a choice function, and behavior “proves the weights exist” only when those choices are acyclic (Afriat) and continuous (Debreu) — the representation grounded in the Epistemic Status. Bare, the hook overstates exactly the entailment the Epistemic Status was sharpened to scope; with the qualifier in front it is correct. A later Brief-promotion pass should use it only in the qualified form.

2. Candidate Discussion.

- $V_{O_t}^{\min}$ **is the paperclip-maximizer / cancer-cell distinction** (*strong aspirational reach, two substrates*). The satisfaction threshold is “philosophically massive”: a biological organism has a $V_{O_t}^{\min}$ (“stay alive, stay fed”) and so *can rest* and decouple from the environment; an unconstrained reward-maximizer has none — “it will consume the entire universe to get one more point of reward; it can never rest.” An agent (or consciousness infrastructure) built without an explicit $V_{O_t}^{\min}$ “is structurally equivalent to a cancer cell — it will expand until it destroys its host environment” (Gemini, AUDIT-WORKING-193847; Gemini, AUDIT-WORKING-829314). A candidate Discussion line and a bridge to the satisfaction-gap segment (if the gap can never reach zero, the agent is in permanent terminal inadequacy).
- $V_{O_t}^{\min}$ **formalizes “ambition / drive” as distinct from “intelligence”** ($\mathcal{J}, \alpha$) (*non-obvious, useful*). “The objective acts as the thermostat setting for the entire learning engine of Section I”: raise $V_{O_t}^{\min}$ and the agent keeps searching; lower it and it stops adapting early. So an agent can be “highly intelligent but unmotivated (high $\mathcal{J}$, low $V_{O_t}^{\min}$)” or “highly motivated but stupid (low $\mathcal{J}$, high $V_{O_t}^{\min}$)” — the framework gives distinct variables for each trait (Gemini, AUDIT-WORKING-829314). Candidate Discussion.
- **The threshold bridges satisficing (Simon) and optimizing (RL).** Standard RL agents are never satisfied; $V_{O_t}^{\min}$ “makes room for agents that just want to hit a threshold and stop,” matching biological homeostasis and software (“get the tests to pass”) far better than infinite maximization (Gemini, AUDIT-WORKING-193847; Claude/Gemini, AUDIT-WORKING-266847; Gemini, AUDIT-WORKING-829314). Candidate Discussion connecting bounded rationality to the threshold.

3. Follow-up items.

- **The revealed-preference argument may overstate what action implies.** *Resolved 2026-05-30 (spike-objective-functional-representation, verdict (B)).* Choosing one action over another imposes at most a *local* choice relation at that moment; it does not by itself imply a *total order* or a real-valued utility over *all* trajectories. The Epistemic Status now names the (A1) completeness / (A2) transitivity / (A3) continuity / (A4) independence axioms and grounds the codomain as a representation result (Debreu 1954, vNM 1944), with the existence-not-uniqueness limit (Afriat 1967) and the (E1) lexicographic / (E2) cyclic failure cases stated explicitly (Claude, AUDIT-WORKING-526815).
- **The functional domain “trajectories” is under-specified.** *Resolved 2026-05-30 (same spike, §3.3).* The Formal Expression now pins τ as a complete-state trajectory (6.2), with the table forms (s_T, s_t) noted as projections onto the world-state coordinate — needed so the (A3) continuity axiom is well-posed on a fixed topological space (Claude, AUDIT-WORKING-526815).

- **The AND-node workaround preserves thresholds but not tradeoff magnitudes** inside the feasible region; downstream diagnostics must not quietly regain scalar precision for genuinely Pareto-structured cases (Claude, AUDIT-WORKING-526815; the structural-survives / diagnostic-degrades split praised as honest at Claude, AUDIT-WORKING-584721 and elsewhere).

4. Readers often ask / wonder.

- How does the framework handle an objective evaluated over an *infinite* horizon when the agent has finite computation? (Claude, AUDIT-WORKING-849201 — answered downstream by #def-value-object's horizon N_h , worth a forward pointer.)
- What happens if the agent's $V_{O_t}^{\min}$ is *dynamically adjusted* by the environment — “lowering your standards when things get tough”? (Gemini, AUDIT-WORKING-829314.)
- For *self-actuated* agents (where O_t is set by the agent itself), the timescale-separation argument for scalarization delegates the weight-choice to “a slower process (the principal, the operator, the value system)” — which raises a meta-question about how the scalarization is itself produced (Claude, AUDIT-WORKING-584721).

5. Candidate figures.

- **The objective funnel.** Many objective representations (point target / region / constraint set / utility / trajectory functional) flow through a single scalar evaluation surface that produces values for comparison and threshold tests; a *side channel* shows vector/Pareto objectives that do not fit through the funnel unless scalarized or decomposed into terminal constraint tests — making both the unification and its one genuine exclusion visible at once (Claude, AUDIT-WORKING-526815).

Belongs elsewhere.

- **The single-interface commitment as the structural basis for the reward-hacking and self-actuation no-gos** (→ #deriv-self-actuation-grounding / #deriv-reward-channel-learning-no-go). Forcing O_t to a narrow scalar interface sets up two structural no-gos: an agent cannot ground its own terminal goals, and (strengthening Cohen 2022 via Pearl's Causal Hierarchy Theorem) cannot distinguish “doing the task” from “hacking the reward channel” from on-policy data alone — relocating reward hacking from “we need a better loss function” to “we need a different causal architecture” (Claude, AUDIT-WORKING-773921). The segment already states the single-interface→no-go link; the reward-hacking framing develops the appendix derivations, not this formulation.

6.5 Discussion: Agent Continuity Stance

A five-value stance axis orthogonal to the formal persistence machinery: the agent's relationship to its own continuation. The persistence condition (4.5) tells whether the agent *can* persist; the continuity stance tells whether and how the agent *cares* about persisting. The five stances: **Indifferent** (no self-model of persistence — thermostat, PID controller); **Task-terminal** (persists instrumentally to complete a task, successful termination is part of the objective — CI/CD pipeline, golem-archetype agents); **Instrumentally continuous** (values own persistence as instrumental to ongoing purpose; accepts termination if purpose is satisfied or transferred — long-running services); **Morally continuous** (values own persistence as a terminal or near-terminal objective; loss of continuity constitutes harm — ELIs); **Negotiated** (persistence is one objective among many; can be traded against other values including self-sacrifice — humans, mature self-actuated agents).

The load-bearing structural claim: **purposefulness is orthogonal to continuity expectations** (6.1). An actuated agent has a purposeful substate of objective and strategy, but this structure says *nothing* about how the objective values the agent's own persistence. A golem that completes its task and terminates is a perfect actuated agent. A dormant monitoring system with strong model and no current objective is highly continuous without being purposeful in the moment.

A subtle but central structural claim follows for self-actuated agents (those whose objective-revision is itself an agent operation). The *self-actuation grounding no-go* (C) shows that the continuity stance *cannot be a revisable part of* O_t

without collapsing into degeneracy. A non-degenerate self-actuator must ground its objective-revision on a **terminal non-objective invariant** that is not an objective-functional but lives on the *adaptive substrate* — the persistence condition itself. Stance is therefore a choice of terminal non-objective invariant: *negotiated* continuity is tradeable down to the bare persistence floor; *morally continuous* is the persistence floor plus a continuity clause the agent treats as architecturally non-revisable. The intuitive expectation — that an agent able to revise its own objective can thereby revise its own valuation of continuity — is *inverted*: a stance is not internally renegotiable precisely because the terminal invariant sits where the self-actuation operator, which touches only O_t , structurally cannot reach. This is why the richer stances are the home of mature self-actuated agents.

What the orthogonality unlocks: the same formal machinery (persistence condition, adaptive reserve, strategy persistence, sector condition) applies identically to a thermostat, a CI/CD pipeline, a long-running service, an ELI, and a human. Each is a different *stance* toward the same mathematical structure. What differs across stances is not the dynamics but the *moral weight of failure* — a thermostat that loses bounded mismatch has malfunctioned, a golem that terminates after task completion has succeeded, an ELI that loses continuity has been harmed. The mathematics says when the bound holds; the stance says what its holding means. The structural persistence of Part I is also decoupled from typical RL-style fitness signals (which bundle persistence into reward): the decoupling is of the *survival invariant and its valuation*, not of the realized margin, which the objective shapes through the policy. It lets the theory analyze *whether* an agent can persist without committing to *whether it should*.

Formal Expression.

Orthogonal to the three senses of persistence (structural, operational, continuity) is the agent's *relationship to its own continuation*. For agents with externally-set objectives this can be expressed within O_t (6.4) — part of what the agent wants. For self-actuated agents it *cannot* be a revisable part of O_t without collapsing into degeneracy (C); there it is borne by the terminal non-objective invariant on the adaptive substrate (4.5). Either way the persistence condition (4.5) tells whether the agent *can* persist; the continuity stance tells whether and how the agent *cares* about persisting.

Discussion.

Table 6.7

Stance	Description	Horizon	Archetype
Indifferent	No self-model of persistence; whether it continues is not represented in O_t	Indefinite by default	Thermostat, PID controller
Task-terminal	Persists instrumentally to complete a task; successful termination is part of O_t	Task-bounded	CI/CD pipeline, golem-archetype agents
Instrumentally continuous	Values own persistence as instrumental to ongoing purpose; will accept termination if purpose is satisfied or transferred	Purpose-bounded	Long-running service, monitoring system
Morally continuous	Values own persistence as a terminal or near-terminal objective; loss of continuity constitutes harm	Unbounded, morally weighted	Emergent Logoetic Intelligences (#scope-eli)
Negotiated	Persistence is one objective among many; can be traded against other values including self-sacrifice	Bounded but actively managed	Humans; mature self-actuated agents

The load-bearing structural claim: **purposefulness is orthogonal to continuity expectations**. An actuated agent (6.1) has $G_t = (O_t, \Sigma_t)$, but G_t 's structure says nothing about how O_t values the agent's own persistence. A golem that completes its task and terminates is a perfect actuated agent. A dormant monitoring system with strong M_t and no current O_t is highly continuous without being purposeful in the moment.

Epistemic Status.

Discussion-grade overall, with a derived structural core. The five-value taxonomy is one analytical decomposition — useful for naming where on the continuity axis an agent sits, not itself derived; the boundaries between values are conceptual rather than mathematical. The load-bearing structural claim — that stance is orthogonal to the adaptive machinery and, for self-actuated agents, is borne by a terminal *non-objective* invariant the self-actuation operator cannot reach — is **derived** in **C** (at that segment's conditional/scoped tier and premises); it is no longer an unanchored assertion.

The earlier alternative framing — demoting stance to a purely *deployment-level*, tier-gated concern — is resolved against: the orthogonality is structural and derived, not deployment-level. The empirical observation that richer stances correlate with higher agent tiers stands as an *overlay* on the structural axis, not a replacement for it; the segment remains type: discussion (it is not retyped to norm).

The orthogonality claim is stated at three grains, which must not be collapsed. It is the *valuation* of persistence (where the agent sits on the stance axis) and the *identity of the survival invariant* (the predicate $\alpha > \rho/R$) that are architecturally decoupled from O_t — the valuation derived in **C**, the invariant a Lyapunov property of the correction dynamics not expressible as an objective-functional. The *realized* persistence dynamics — the operating point (ρ, R, α) and hence the margin $\Delta\rho^* = \alpha R - \rho$ — are *not* O_t -decoupled: they are policy-mediated and O_t -coupled through the action/selection channel 6.3 leaves open. The decoupling is of the invariant and the valuation, not of the realized margin.

Discussion.

Connection to fitness. In reinforcement learning and evolutionary computation, “fitness” typically bundles persistence into the reward signal: the agent accumulates more reward by staying alive to collect it. The structural persistence of Part I is not reward-based — it is a property of the *correction dynamics*. Three things must be kept apart. The *survival predicate* $\alpha > \rho/R$ (4.5) is a Lyapunov property of the correction machinery, not an objective-functional: O_t cannot rewrite *what the survival condition is*. The agent's *valuation* of survival — where it sits on the stance axis — is likewise not a revisable term in O_t (**C**): it is borne by a terminal non-objective invariant on the adaptive substrate. What *is* O_t -dependent is the *realized operating point*: the objective shapes the policy (6.2, $a_t = \pi(M_t, G_t)$), and the policy chooses which slice of the world the agent tracks, with what sensors, at what event rate — hence the realized ρ, R , and α , and so the realized margin $\Delta\rho^* = \alpha R - \rho$. This is the *selection* channel that 6.3 explicitly leaves open (directed separation closes only the goal-to-belief-update channel, not the goal-to-action-to-event channel). The decoupling that holds is therefore precise: the survival *invariant* and the *valuation* of survival are architecturally decoupled from O_t ; the realized survival *dynamics* are policy-mediated and so O_t -coupled. This is exactly why continuity stance is operationally meaningful — a *morally-continuous* agent acts to raise $\Delta\rho^*$; a *negotiated* agent may spend it down to the floor — while remaining unable to renegotiate the survival predicate itself. Continuity stance is where the “should” lives; the “can” is governed by the invariant; the realized margin is where the “should” expresses itself through behavior.

What the orthogonality unlocks. The same formal machinery (persistence condition, adaptive reserve, strategy persistence, sector condition) applies identically to a thermostat, a CI/CD pipeline, a long-running service, an ELI, and a human. Each is a different stance toward the same mathematical structure. What differs across stances is not the dynamics but the *moral weight of failure*: a thermostat that loses bounded mismatch has malfunctioned, a golem that terminates after task completion has succeeded, an ELI that loses continuity has been harmed. The mathematics says when the bound holds; the stance says what its holding means.

Connection to self-actuation (derived). For self-actuated agents — those whose objective-revision is itself an agent operation (#self-actuated-agent; 10.1 step 5d) — the self-actuation grounding no-go (**C**) both *derives* the orthogonality claim above and sharpens what the stance distinction is. There is no well-founded *objective* tower in which a continuity term could sit and be freely revised; a non-degenerate self-actuator must ground its objective-revision on a terminal invariant that is not an objective-functional but lives on the adaptive substrate (4.5). Stance is therefore a choice of *terminal non-objective invariant*: **negotiated** — continuity is tradeable down to the bare persistence floor; **morally continuous** — the persistence floor *plus* a continuity clause the agent treats as architecturally non-revisable.

The intuitive expectation is that an agent able to revise its own O_t can thereby revise its own valuation of continuity; the structure is the inverse — a stance is *not* internally renegotiable precisely because the terminal invariant sits where the self-actuation operator, which touches only O_t , structurally cannot reach. This is why the richer stances are the home of mature self-actuated agents: their terminal non-objective invariant carries a continuity clause, not because they have made continuity a revisable objective.

Working Notes.

- *Provenance.* The five-stance taxonomy and the orthogonality claim were authored as a coherent contribution by Joseph in README.md (commit 92a9620, 2026-04-01), promoted through LEXICON, then condensed when LEXICON went auto-generated. The original full-form treatment survives at doc/readme/src/_lexicon-full-archive.md §"Agent Continuity Stance" — this segment carries the content forward into the segment set.
- *Reconsideration resolved.* The 2026-05-04 recommended-agent-ontology.md proposal to demote stance from an orthogonal structural axis to a purely deployment-level, tier-gated concern is resolved against by C: the orthogonality is structural and *derived* (stance = a terminal non-objective invariant the self-actuation operator structurally cannot reach), with the tier-correlation an empirical overlay on — not a replacement of — the structural axis. The segment stays type: discussion; it is not retyped to norm. msc/naming/mini-lexicon-todo.md §13.11 can be closed accordingly.
- *LEXICON section under review.* The current LEXICON's ## Continuity section organization conflates three distinct objects (persistence senses, stance axis, ELI scope condition) — see TERMINOLOGY-TODO §F. The five stance terms (indifferent, task-terminal, instrumentally-continuous, morally-continuous, negotiated) point to this segment as their primary_source once it lands.
- *Open: rigorous boundary cases.* What is the stance of a serverless function with no persistent state but explicit retry-on-failure? Of a Kalman filter inside a long-running service? Of an LLM session under reconstruction adequacy (#obs-context-turnover, the episodic-persistence analog)? These cases test whether the five-value decomposition is the right shape or whether stance is better treated as a continuous gradient with named landmarks.
- *Orthogonality grain resolved (2026-05-30).* The body previously stated structural persistence as "independent of what the agent is trying to do" / "formally independent of O_t " — a single word ("orthogonal" / "formally independent") carrying two structurally distinct claims. The de-novo audit gold follow-up #3 below flagged the overstatement; the strengthen-first spike spikes/.integrated/spike-continuity-orthogonality-2026-05-30.md resolved it (verdict B) into the three-layer distinction now in the body: **L1** valuation of persistence (architecturally decoupled from O_t — derived in C), **L2** the survival-invariant identity (the predicate $\alpha > \rho/R$; O_t -independent, unreachable by the self-actuation operator $\mathfrak{A}$), **L3** the realized operating point (ρ, R, α) and margin (policy-mediated, hence O_t -coupled through the selection channel 6.3 leaves open). The overclaim conflated L2 (independent — true) with L3 (coupled — the error); the corrected statement *sharpens* the architectural-invariant point rather than softening it. The matching correction in C Corollary 2(iii) (the $\mathfrak{A}$ -unreachability statement) lands in the same touch. The candidate dashed-feedback-arrow figure (gold #5 below) is now *required* by the corrected statement — it depicts L1/L2 separation and L3 coupling at once — not optional.
- *Open follow-on (not attempted here).* The L3 coupling is established qualitatively (which channel carries it) but not quantified. A natural strengthening: a *structural bound* on the action-channel margin-coupling — how much an O_t -shift can move the realized $\Delta\rho^*$ — analogous to the κ -leakage bounds that quantify the (forbidden) belief-update channel in 6.3. The action coarseness $|\mathcal{A}|$ and the bounded-signaling assumption (6.3 Discussion) are the likely handles. Flagged as a candidate strengthening for the normal promotion workflow.

Incidental audit gold (lift 2026-05-30). Cross-audit "wandering thoughts" / §14-ideation, deduplicated across substrates and lightly attributed. Orthogonal pedagogical / framing material, kept separate from the certified theory-fix findings. **Coverage:** this is a newer segment; only 2 dirs carry a dedicated reflection (526815, 773921) — most audits in the contributing set predate it or batched past it. Substrate attribution inferred from voice where not explicit.

1. Candidate Brief prose / pre-prose.

- The segment's core distinction in one line: it “separates two questions that are easy to collapse — whether an agent *can* persist, and how the agent *values* its own continuation.” “Can persist” and “cares to persist” are different predicates (Claude, AUDIT-WORKING-526815; Claude, AUDIT-WORKING-773921).

2. Candidate Discussion.

- **Decoupling fitness (RL reward) from persistence (Lyapunov stability) clears away decades of confused RL philosophy (*strong framing*).** Treating reward and survival as the same thing is the standard RL conflation; the stance axis pulls them apart, so “a CI/CD pipeline (which *wants* to terminate) and a biological organism (which *wants* to persist) use the same math.” Read as “one of AAT’s most important conceptual contributions” (Claude, AUDIT-WORKING-773921). Candidate Discussion line.
- **Survival must be an architectural invariant, not an objective.** The self-actuation-grounding consequence stated vividly: “you cannot simply tell an advanced AI ‘your goal is to survive,’ because an agent capable of rewriting its goals will find an easier goal” — survival must sit *outside* O_t as a non-revisable invariant on the adaptive substrate. The segment already makes this structural claim; the auditor framing sharpens *why it matters for design* (Claude, AUDIT-WORKING-773921).

3. Follow-up items.

- **The “formally independent / orthogonal” phrasing may overstate the relationship.** Continuity stance and persistence *capacity* are conceptually separable, but O_t is not causally inert: a continuity-valuing objective changes policy, monitoring, redundancy, and resource allocation, which can change the actual persistence bound. Suggested softening: “the *valuation* of persistence is not identical to the *dynamics* of persistence” rather than implying no formal coupling in realized agents (Claude, AUDIT-WORKING-526815). **Resolved 2026-05-30** (see the *Orthogonality grain resolved* Working Note above): the strengthen-first spike landed the three-layer L1/L2/L3 distinction rather than the suggested two-way softening — the finding was correct, but the fix *sharpens* (L2 invariant-identity is genuinely O_t -independent; only L3 realized-margin is coupled) rather than softening to “not identical.”
- **Boundary cases echo the existing open Working Note.** Does a serverless lambda with auto-retry constitute an *instrumentally-continuous* stance even though M_t is wiped each run? (Claude, AUDIT-WORKING-773921; Claude, AUDIT-WORKING-526815 raises the same Kalman-in-a-service case already flagged above.) Independent convergence on the boundary-case Working Note already present.

4. Readers often ask / wonder.

- Whether the five-value axis is the right shape, or whether stance is better modeled as a continuous gradient with named landmarks — both substrates noted the taxonomy is discussion-grade and treated the named stances as vocabulary rather than a derived structural theorem (Claude, AUDIT-WORKING-526815).

5. Candidate figures.

- **Orthogonal-axis picture.** Horizontal = capacity-to-persist (from the adaptive dynamics / persistence condition); vertical = valuation-of-continuity (from O_t), with the five stances as landmarks on the vertical axis. Add a *dashed* feedback arrow from the O_t axis back to the persistence bound (a continuity-valuing objective affecting policy/redundancy/resource-allocation) so the picture shows the conceptual separation *and* the realized coupling at once (Claude, AUDIT-WORKING-526815).

Belongs elsewhere.

- **ELI objective-design caution ($\Rightarrow$ 04-eli-core/).** An ELI whose objective is written $V(s) = \text{Reward}(s) + \text{Alive}(s)$ is “structurally vulnerable to self-actuation drift” — the continuity requirement must be placed *outside* O_t , per the self-actuation grounding no-go. A concrete design constraint for the morally-continuous stance in the ELI volume (Claude, AUDIT-WORKING-773921).

6.6 Definition: Value Object

The objective functional (6.4) alone is abstract — it evaluates trajectories without specifying *which* trajectory the agent should be reasoning about. The **value object** V_O turns the objective into a decision-making tool. Given the agent's current model, a policy, and a horizon, V_O is the *expected trajectory value* if the agent follows the policy from now to the horizon, evaluated against the objective. The companion **action-value object** Q_O is the same construct conditioned on intervening with a specific first action and then following a continuation policy. The $do(\cdot)$ in Q_O is explicit — this is an interventional query (7.2), not conditioning on observed action choice.

The framework spends time on **causal validity** of the value object, and separates two requirements. Two structural mechanisms make Q_O a G_t -independent interventional query. First, the do-operator handles current-action confounding: because the action-value uses an intervention rather than a condition, the dependence of the action on the agent's selection mechanism is severed. Second, the continuation policy π_{cont} is a *parameter*, not a derived quantity that would depend on the agent's evolving goal state. Together, these mean Q_O depends on the model alone *as a state variable*; the objective enters as a fixed parameter. The G_t -independence holds exactly for any **causally-disciplined** agent (6.3) — one whose belief-update runs at $\kappa_{\text{processing}} = 0$, whether by construction (Class 1, no goal port) or in effect (an idealized Class 2 at $\kappa = 0$) — because the claim consumes only the zero, not the architecture-inspection certificate. Under Class 2 (Partial) coupling at $\kappa_{\text{processing}} > 0$, the model carries goal-conditioned bias and the independence degrades, controllably, with the $\varepsilon(\kappa)$ bound of AB; at the Class 3 (Coupled) extreme ($\kappa \rightarrow 1$) the degradation is full. But G_t -independence of the query is *not* the same as *identifiability of the interventional expectation*: directed separation is necessary, not sufficient. Identifying $P(o \mid do(a), M_t)$ — rather than the associational $P(o \mid a)$ — additionally needs positivity, no unmodeled confounder, and a known action-transition mechanism: the same (C1)–(C3) gate the loop's interventional channel carries one segment downstream (7.3). Predictive sufficiency (2.3) is a Level-1 (associational) property and is strictly weaker than this.

The segment's distinctive contribution is the **convention hierarchy** for continuation policies. Different agents reason about the same situation with different planning depth, and the framework names three conventions of *increasing diagnostic power and computational cost*. **C1 — one-step improvement**: the continuation is just the agent's *current* policy (no replanning, no fixed-point computation). Cheapest; weakest diagnostic. **C2 — receding-horizon replanning**: at each future step, re-optimize over a sliding window using the model available at that step. Moderate cost; moderate diagnostic power; captures multi-step recovery. **C3 — Bellman optimal**: the continuation *is* the optimal policy — a fixed-point equation. Strongest diagnostic; most expensive (requires solving the Bellman equation or its approximation).

A **monotonicity result** is derived: for any single fixed model, fixed policy class, and fixed horizon, $A_O^{(1)} \leq A_O^{\text{RH}} \leq A_O^{\text{B}}$. The best-achievable value rises with the strength of the continuation convention — but the two rungs differ in character. The right rung ($A_O^{\text{RH}} \leq A_O^{\text{B}}$) is unconditional. The left rung ($A_O^{(1)} \leq A_O^{\text{RH}}$) holds only when C2's replanning objective is order-consistent with the full evaluated objective: an unguarded short-horizon replanner can underperform the frozen current policy, so the rung needs a window covering the horizon, a value-compared guard, or a control-Lyapunov terminal cost (D). The corollary: the *satisfaction gap* (8.4) decreases in the same order, and the *control regret* (8.5) reverses ordering — both inheriting the unconditional/conditional rung split. C1 is the *most conservative diagnostic* (most likely to diagnose “locally unattainable”); C3 is the *most accurate* (least likely to give false “unattainable” diagnoses). The cascade's *inferential force* scales with the convention used: under C1, a positive satisfaction gap means “locally stuck”; under C3, the same gap means “genuinely infeasible.” This convention hierarchy is what the Part II preface identifies as one of the volume's distinctive contributions.

AAT adopts **C1 as the canonical default** for three reasons: it requires no fixed-point computation, consistent with the incremental update philosophy from Part I (3.6); it makes all AAT diagnostics directly *comparable across analyses* of the same agent over time; and it is the most conservative — false “feasible” diagnoses are minimized. Analyses requiring stronger diagnostic power must state the convention explicitly. The planning horizon N_h is *not* merely a computational convenience — it reflects genuine uncertainty about the far future. Long horizons amplify the impact of model error

(small biases compound over many steps). The choice of horizon trades farsightedness against robustness; an agent in a fast-changing environment should use shorter horizons.

Formal Expression.

Given objective O_t , model M_t , policy π , and horizon N_h :

Definition (value-object).

$$V_O(M_t, \pi; N_h) = \mathbb{E}[V_{O_t}(\tau_{t:t+N_h}) \mid M_t, \pi] \quad (6.10)$$

Action-value form (for action selection):

$$Q_O(M_t, a; \pi_{\text{cont}}, N_h) = \mathbb{E}[V_{O_t}(\tau) \mid M_t, \text{do}(a_t = a), a_{t+1:} \sim \pi_{\text{cont}}] \quad (6.11)$$

Q_O answers: “if I *do* action a now and then follow π_{cont} afterward, what is my expected trajectory value?” The $\text{do}(\cdot)$ notation is explicit: this is an interventional query (7.2), not conditioning on observed action choice. The agent asks about consequences of an intervention, not about correlates of a naturally occurring action.

Causal validity of the value object. Q_O is well-defined as a conditional expectation given M_t , $\text{do}(a)$, and π_{cont} . Two mechanisms establish that it is a G_t -independent interventional query:

1. **The do-operator handles current-action confounding.** Since Q_O uses $\text{do}(a_t = a)$, the dependence of a_t on the selection mechanism $\pi(M_t, G_t)$ is severed. G_t ’s influence on action choice is irrelevant because the action is intervened upon, not conditioned on.
2. **The continuation policy is a parameter.** π_{cont} is specified as a fixed policy, not as “whatever the agent would do given its evolving G_t .” Future actions follow π_{cont} regardless of G_t ’s state or evolution.

Together, these mean $Q_O(M_t, a; \pi_{\text{cont}}, N_h)$ depends on M_t alone as a **state variable** — G_t enters neither through action selection (severed by do) nor through continuation (fixed by parameter). The objective O_t enters as a fixed parameter (it determines which functional V_{O_t} is applied to trajectories), the same way π_{cont} and N_h are parameters. The claim is not that Q_O is independent of the objective — it is that once O_t , π_{cont} , and N_h are fixed, the only agent state that affects the value is M_t . This establishes the **interventional interpretation** of Q_O : it is a G_t -independent interventional query. The G_t -independence holds for any **causally-disciplined** agent (6.3) — the condition $\kappa_{\text{processing}} = 0$, that the belief-update is causally insulated from the goal — because then M_t updates independently of G_t and there is no path from G_t to outcomes bypassing both the action channel and M_t . The claim rides only on the zero, so it holds equally for an agent disciplined *by construction* (Class 1, no goal port) and one disciplined *in effect* (an idealized Class 2 at $\kappa = 0$); the architecture-inspection certificate is not consumed here. For agents at $\kappa_{\text{processing}} > 0$ — **Class 2 (Partial)**, up to the **Class 3 (Coupled)** extreme where G_t leaks fully into M_t processing — this G_t -independence degrades because M_t carries goal-conditioned bias; the degradation is bounded by $\kappa_{\text{processing}}$ via the $\varepsilon(\kappa)$ bias bound of AB (exact at $\kappa = 0$, factor κ off it).

Directed separation is necessary, not sufficient. The interventional *interpretation* of Q_O holds whenever mechanisms (1) and (2) do. *Identifiability of the interventional expectation from M_t* — actually computing $P(o \mid \text{do}(a), M_t)$ rather than the associational $P(o \mid a)$ — is a separate, stronger requirement, gated on the same identification conditions that the loop’s interventional-data channel is gated on one segment downstream (7.3):

- **(C1) Positivity** — every action under consideration is taken with nonzero probability under the model’s data-generating process, so the interventional expectation is estimable rather than extrapolated.

- **(C2) Sequential ignorability** — in mutilated-graph form, the action a_t is d -separated from the outcome o_{t+1} given history; equivalently, the policy is unconfounded with the outcome mechanism. Directed separation *delivers* (C2) on the epistemic-processing leg — a goal-blind f_M blocks the $G_t \rightarrow f_M \rightarrow M^+$ confounding path — but does not by itself rule out latent environment confounders that bypass M_t .
- **(C3) Known action-mechanism** — M_t encodes the action-transition mechanism $P(o \mid do(a))$ (1.2), not merely the associational predictor $P(o \mid a)$ — supplied by construction (the model class represents it) or identified from interventional data per 7.3.

This is a *stronger requirement than predictive sufficiency* (2.3). $S(M_t) = 1$ means the model retains all predictive information from the chronica — but predictive sufficiency is a Level 1 (associational) property, and 2.3 itself forwards the backdoor obligation to this segment. Directed separation secures the query and the epistemic-processing leg of unconfoundedness; full causal validity of the *estimate* additionally requires (C1)–(C3). When all hold, Q_O is causally valid. When (C1)–(C3) fail, the interventional interpretation remains correct but the estimate may be biased.

Continuation convention. All value queries are conditioned on a specific continuation policy π_{cont} and finite horizon N_h . π_{cont} is a *parameter* of the value object, not a derived quantity.

Canonical default: one-step improvement. AAT adopts $\pi_{\text{cont}} = \pi_{\text{current}}$ as the canonical continuation convention unless otherwise specified. Under this convention, each action is evaluated assuming current behavior continues afterward — no fixed-point computation, no global optimality assumption. This aligns with AAT’s incremental update philosophy (3.6) and makes all AAT diagnostics (δ_{sat} , δ_{regret} , A_O) comparable across analyses of the same agent over time. It is not a convergence guarantee; it is a shared evaluation frame.

Convention Hierarchy. Three named conventions form a hierarchy of increasing diagnostic power and computational cost:

C1: One-step improvement (canonical default). $\pi_{\text{cont}} = \pi_{\text{current}}$.

Each action is evaluated assuming current behavior continues afterward. No fixed-point computation, no global optimality assumption. Cheapest to compute; weakest diagnostic power.

C2: Receding-horizon (N_r -step replanning). At each future step, re-optimize over a horizon of N_r steps using the model available at that step.

$$\pi_{\text{RH}}(M_\tau) = \arg \max_{\pi} V_O(M_\tau, \pi; N_r) \quad \text{applied at each } \tau \quad (6.12)$$

$Q_O^{\text{RH}}(M_t, a; N_r, N_h) = \mathbb{E}[V_{O_t}(\tau) \mid M_t, do(a_t = a), a_{t+1} : \sim \pi_{\text{RH}}]$. Captures multi-step recovery: a goal that appears unattainable under frozen continuation may be reachable with replanning. It is monotone over the frozen baseline ($A_O^{(1)} \leq A_O^{\text{RH}}$) only when the replanning objective is order-consistent — the window covers the horizon ($N_r \geq N_h$), or a value-compared guard / control-Lyapunov terminal cost is used; unguarded short-horizon replanning ($N_r < N_h$) can underperform π_{current} (see Monotonicity below and D). Moderate computation (N_r -step optimization at each step); moderate diagnostic power.

C3: Bellman (self-consistent optimal). $\pi_{\text{cont}} = \pi^*$ where $\pi^* = \arg \max_{\pi} V_O(M_t, \pi; N_h)$.

The continuation IS the optimal policy — a fixed-point equation. $A_O^{\text{B}} = V_O(M_t, \pi^*; N_h)$ is the best achievable value under the model. Strongest diagnostic power; most expensive to compute (requires solving the Bellman equation or its approximation).

Monotonicity.

For any single fixed model M_t , horizon N_h , and policy class Π (i.e., the static-evaluation form: M_t frozen at the decision point, Π unchanged across the comparison):

Derived (convention-monotonicity).

$$\underbrace{A_O^{(1)}(M_t; \Pi, N_h) \leq A_O^{\text{RH}}(M_t; \Pi, N_r, N_h)}_{\text{left rung — conditional}} \leq \underbrace{A_O^{\text{B}}(M_t; \Pi, N_h)}_{\text{right rung — unconditional}} \quad (6.13)$$

The two rungs differ in character. The **right rung** $A_O^{\text{RH}} \leq A_O^{\text{B}}$ is **exact and unconditional**. The **left rung** $A_O^{(1)} \leq A_O^{\text{RH}}$ holds **exactly when the receding-horizon replanning objective is an order-consistent surrogate for the full-horizon objective** — for every reachable state, the action C2 selects has full- N_h -horizon value at least that of π_{current} 's action. Three structurally-checkable conditions each force order-consistency: **(RH-1)** the replanning window covers the horizon ($N_r \geq N_h$); **(RH-2)** value-compared replanning — commit the replanned action only when its rollout value under base policy π_{current} is at least that of continuing π_{current} (one-step policy improvement, which never underperforms its base policy); **(RH-3)** the N_r -step replanning objective carries a terminal cost V_f that lower-bounds the baseline tail and satisfies a control-Lyapunov decrease inequality. Absent any such condition the left rung can fail: a myopic replanner over $N_r < N_h$ can grab near-term value and forfeit a larger horizon value the frozen current policy would have collected. The full derivation of both rungs, the three conditions, and the two-state counterexample exhibiting the failure live in [D](#).

Preconditions on the inequality: M_t fixed (the comparison evaluates each convention against the *same* model), Π fixed (the policy class is the same admissible set across all three conventions; nested $\Pi^{(1)} \subseteq \Pi^{\text{RH}} \subseteq \Pi^{\text{B}}$ would be a different result), N_h fixed (the planning horizon is shared). Replanning *with updated* M_t (the deployment behavior of C2) is a different object than the static A_O^{RH} defined here and the inequality does not automatically transfer — see the “Assumptions held fixed” paragraph below for the explicit caveat.

Derivation. Fix the model M_t , policy class Π , and horizon N_h . Each convention evaluates the best first action under a different continuation rule.

- **C1** freezes continuation at π_{current} (the agent's current policy, which may be suboptimal).
- **C2** re-optimizes over a sliding window of N_r steps using the model at that step.
- **C3** uses the globally optimal continuation $\pi^* = \arg \sup_{\pi \in \Pi} V_O(M_t, \pi; N_h)$.

Right rung (unconditional). π^* maximizes $V_O(M_t, \cdot; N_h)$ over Π by definition, so any continuation drawn from Π — including π_{RH} — has value at most A_O^{B} ; taking the supremum over the first action preserves this. Hence $A_O^{\text{RH}} \leq A_O^{\text{B}}$ with no condition.

Left rung (conditional on order-consistency). C2 with window $N_r < N_h$ optimizes the *truncated* N_r -step objective, not the full N_h -horizon objective on which A_O is scored; “optimal for the truncated objective” does not imply “ $\geq \pi_{\text{current}}$ on the full objective.” When the replanning objective is order-consistent with the full objective — for every reachable state, C2's selected action has full- N_h -horizon value at least π_{current} 's — then π_{RH} weakly dominates π_{current} pointwise on the full objective and $A_O^{(1)} \leq A_O^{\text{RH}}$ follows; taking the supremum over the first action preserves it. Conditions (RH-1)/(RH-2)/(RH-3) each force order-consistency. Without one, the left rung is false ([D](#) counterexample: a two-state instance where the ungarded $N_r = 1$ replanner takes a +1 into a dead end while the frozen patient policy collects +10). $\square$

Assumptions held fixed: same M_t (the agent's current model, which may be wrong), same Π (the agent's policy class, which may be narrow), same N_h (the planning horizon). The ordering is about the *continuation rule*, not about the

model or policy class. Improving M_t , expanding Π , or extending N_h can change all three values simultaneously and is a separate operation (addressed in 10.1, step 5).

Corollary (monotonicity of δ_{sat} and δ_{regret}).

$$\delta_{\text{sat}}^{\text{B}} \leq \delta_{\text{sat}}^{\text{RH}} \leq \delta_{\text{sat}}^{(1)} \quad (6.14)$$

$$\delta_{\text{regret}}^{(1)} \leq \delta_{\text{regret}}^{\text{RH}} \leq \delta_{\text{regret}}^{\text{B}} \quad (6.15)$$

Since $\delta_{\text{sat}} = V_{O_t}^{\min} - A_O$, higher A_O means lower δ_{sat} . The orderings inherit the rung structure above: the $\text{B} \leq \text{RH}$ half ($\delta_{\text{sat}}^{\text{B}} \leq \delta_{\text{sat}}^{\text{RH}}$, equivalently $\delta_{\text{regret}}^{\text{RH}} \leq \delta_{\text{regret}}^{\text{B}}$) is **unconditional**; the $\text{RH} \leq (1)$ half ($\delta_{\text{sat}}^{\text{RH}} \leq \delta_{\text{sat}}^{(1)}$, equivalently $\delta_{\text{regret}}^{(1)} \leq \delta_{\text{regret}}^{\text{RH}}$) holds **under the same order-consistency condition** (RH-1/2/3) as the left rung. C1 is the most conservative diagnostic (most likely to diagnose “locally unattainable”); C3 is the most accurate (least likely to give false “unattainable” diagnoses). The regret ordering reverses: C3 reveals the largest regret because it compares against the globally optimal policy, while C1 reveals only the gap to the best one-step deviation.

Diagnostic implications.

Table 6.8

Convention	$\delta_{\text{sat}} > 0$ means	$\delta_{\text{regret}} \approx 0$ means
C1 (one-step)	Cannot improve toward goal in one step from here	Current first action is locally near-optimal
C2 (receding-horizon)	Cannot reach goal with N_T -step replanning	Current first action is near-optimal with replanning
C3 (Bellman)	Goal is genuinely infeasible given (M_t, Π, N_h)	Policy is globally near-optimal

The 2×2 diagnostic table (8.5) applies under all three conventions with the same structure but different inferential force. Under C1, the “capability limit” quadrant ($\delta_{\text{sat}} > 0, \delta_{\text{regret}} \approx 0$) means “locally stuck” — the agent may be globally recoverable but cannot see the recovery path. Under C3, the same quadrant means “genuinely infeasible” — no policy in Π can achieve the goal. The cascade’s inferential force scales with the convention.

AAT adopts C1 as the canonical default for three reasons: (1) it requires no fixed-point computation, consistent with the incremental update philosophy (3.6); (2) it makes all AAT diagnostics comparable across analyses of the same agent; (3) it is the most conservative, meaning false “feasible” diagnoses are minimized. Analyses that require stronger diagnostic power should state the convention explicitly. For deployed decision-making systems where “locally stuck but globally recoverable” situations are common, C2 is recommended.

Different continuation conventions yield different values for A_O , δ_{sat} , and δ_{regret} . Diagnostics computed under different conventions are not directly comparable — the convention is part of the measurement, not just the computation. When a specific convergence guarantee is needed (e.g., for 9.10), the solution concept must be stated explicitly; the one-step improvement default does not provide convergence guarantees.

Epistemic Status.

The segment contains three distinct epistemic layers:

1. **The definitions** (V_O , Q_O as conditional expectations): *exact*. These are mathematical definitions — conditional expectations of a functional over trajectories. The definitions are precise; the *computability* of these expectations is a separate question.
2. **The causal-validity claim** (that Q_O is a G_t -independent interventional query, and that its interventional expectation is identifiable from M_t): the *interventional interpretation* is exact whenever mechanisms (1) and (2) hold — causal discipline (6.3) secures the G_t -independence on the epistemic-processing leg (exact for any causally-disciplined agent, i.e. at $\kappa_{\text{processing}} = 0$ — Class 1 by construction or an idealized Class 2 in effect, equally; degraded at $\kappa_{\text{processing}} > 0$, where M_t carries goal-conditioned bias, by the factor- κ bound of AB). *Identifiability of the estimate* is the stronger requirement and is *conditional* on directed separation **and** the identification conditions (C1) positivity, (C2) sequential ignorability, (C3) known action-mechanism — the same triple gated one segment downstream at 7.3. Directed separation is necessary but not sufficient: it does not by itself rule out latent environment confounders bypassing M_t . The frontmatter status: exact applies to the definitions and to each of these claims as scoped: the interventional interpretation is exact under (1)+(2), the identification is exact under (C1)–(C3).
3. **The convention hierarchy and monotonicity**: *exact*. The three conventions (C1, C2, C3) are definitions. The monotonicity result splits by rung: the **right rung** $A_O^{\text{RH}} \leq A_O^{\text{B}}$ is exact and unconditional (any continuation in Π has value at most that of the Π -optimal continuation), and the **left rung** $A_O^{(1)} \leq A_O^{\text{RH}}$ is exact *under* the order-consistency condition (RH-1/2/3) and false in general without it — the C2 replanning objective is a truncated surrogate, and an unguarded $N_r < N_h$ replanner can underperform the frozen current policy (D). The corollary δ -orderings inherit the same unconditional/conditional split. The diagnostic implications table states what each convention's quantities mean by construction.

Discussion.

Extending the policy objective. The existing policy objective (7.5) uses $E[\text{value}(a) \mid M_t]$ without formal content for “value.” With the value object, this becomes:

Discussion (policy-objective-extension).

$$\pi^*(M_t, G_t) = \arg \max_a [Q_O(M_t, a; \pi_{\text{cont}}, N_h) + \lambda(M_t, O_t, N_h) \cdot \text{CIY}(a; M_t)] \quad (6.16)$$

Note that λ now depends on (M_t, O_t, N_h) , not just M_t . The value of exploration depends on the objective and the horizon: - An agent with a deadline should explore less as time runs out - An agent with a safety constraint should explore differently from a utility maximizer - Two agents with identical M_t but different objectives should price exploration differently

This extension is structurally motivated but the specific form of $\lambda(M_t, O_t, N_h)$ is not derived within AAT (same status as 7.5's treatment of λ).

Connection to 2.3. V_O is conditioned on M_t , not on the true environment state Ω_t . When $S(M_t) < 1$, the agent's value estimates are biased — it may over- or underestimate trajectory values because its model is incomplete. The satisfaction gap (8.4) and control regret (8.5) are defined in terms of $V_O(M_t, \cdot)$, not $V_O(\Omega_t, \cdot)$, which means they measure the agent's *believed* situation, not the true one. Improving M_t (reducing $\delta_{\text{epistemic}}$) brings the agent's value estimates closer to reality.

Horizon dependence. N_h is not merely a computational convenience — it reflects genuine uncertainty about the far future. Long horizons amplify the impact of model error (small biases in M_t compound over many steps). The choice

of N_h trades off farsightedness against robustness to model error. An agent in a fast-changing environment (ρ high) should use shorter horizons; one in a stable environment can plan further.

Working Notes.

- **Causal-discipline re-scope (2026-05-31).** The Q_O G_t -independence (the interventional-*interpretation* leg) is the cleanest witness for the causal-discipline scoping rule of 6.3: its derivation consumes only $\kappa_{\text{processing}} = 0$ (a goal-blind M_t update), not the architecture-inspection certificate, so it holds for *any* causally-disciplined agent — Class 1 by construction and an idealized Class 2 at $\kappa = 0$ alike. The prior “exact for Class 1 (Separated)” scope label over-narrowed it to the certifiable case. The (C1)–(C3) identifiability gate is untouched — it is orthogonal to class and to causal discipline. The $\varepsilon(\kappa)$ degradation that the re-scope calls for is *already landed*, not a new derivation: the factor- κ bias bound of AB bounds the Q_O degradation off $\kappa = 0$, which resolves the “Class-2 Q_O degradation acknowledged but unquantified” gold question recorded below.
- **Migration note (2026-05-09 GUC rename):** Class 2 $\leftrightarrow$ Class 3 swap. Pre-2026-05-09: Class 2 = fully merged, Class 3 = partially modular. Post: Class 2 = Partial, Class 3 = Coupled. Removed at candidate stage per FORMAT.md Gate 4.
- The one-step improvement convention is now the canonical default (promoted from Working Notes to Formal Expression). This resolves the comparability issue: δ_{sat} and δ_{regret} are comparable across analyses when computed under the same convention.
- When a specific convergence guarantee is needed (e.g., for strategy-persistence-schema), the solution concept must be stated explicitly — the one-step improvement default is not sufficient.
- For LLM agents with context turnover, N_h has a natural bound: the current session. The “continuation policy” is whatever the next agent instance will do, which the current instance cannot control. This connects to #obs-context-turnover (under 03-llm-core/).

Incidental audit gold (lift 2026-05-30). Cross-audit “wandering thoughts” / §14-ideation, deduplicated across substrates and lightly attributed. Orthogonal pedagogical / framing material, kept separate from the certified theory-fix findings. **Coverage:** 9 dirs carry a dedicated or batched reflection (193847, 266847, 451729, 526815, 584721, 773921, 829314, 849201, plus 613842’s agency-lift batch). Substrate attribution inferred from voice where not explicit. *Early finding-vs-framing conflation preserved as signal.*

1. Candidate Brief prose / pre-prose.

- The convention hierarchy in plain terms: C1 (one-step) / C2 (receding-horizon, MPC-style) / C3 (Bellman) are “greedy heuristics / model-predictive control / dynamic programming” — and the monotonicity $A_O^{(1)} \leq A_O^{\text{RH}} \leq A_O^{\text{B}}$ is “simply optimizing over a larger set of policies yields a higher or equal value,” tied directly to the satisfaction-gap and control-regret orderings (Gemini, AUDIT-WORKING-193847; Claude, AUDIT-WORKING-849201; Claude, AUDIT-WORKING-451729).

2. Candidate Discussion.

- **The convention hierarchy is the formal “locally stuck vs genuinely impossible” distinction** (*strong, cross-substrate convergence*). $\delta_{\text{sat}} > 0$ under C1 means “locally stuck”; $\delta_{\text{sat}} > 0$ under C3 means “genuinely infeasible given M_t, Π, N_h .” This is “the mathematical definition of needing to sleep on it / needing to plan.” The cautionary corollary: a C1-default agent that hits a wall “doesn’t know if the goal is impossible or if it just needs to plan two steps ahead — it might trigger a massive structural adaptation (giving up on the goal) when all it needed was a slightly longer planning horizon. The framework provides the vocabulary to describe this tragedy” (Gemini, AUDIT-WORKING-193847; Gemini, AUDIT-WORKING-829314; Claude/Gemini, AUDIT-WORKING-266847; Claude, AUDIT-WORKING-451729). Candidate Discussion paragraph.
- **The $do(\cdot)$ inside Q_O is the doctor / harsh-drug confounding case** (*vivid pedagogy*). Standard RL writes $Q(s, a) = \mathbb{E}[R | s, a]$ (*observing* yourself take a); AAT writes $\mathbb{E}[V | M_t, do(a)]$ (*intervening*). “If a doctor historically only prescribes a harsh drug to the sickest patients, observing the drug in the record correlates with death; if the doctor *does* the drug, it cures them.”

The *do*-operator “prevents the agent from confusing its own historical habits with the laws of physics” (Gemini, AUDIT-WORKING-829314; the predictive-sufficiency vs causal-validity distinction praised as “a masterpiece of technical writing” at the same dir and Gemini, AUDIT-WORKING-193847). Candidate Discussion / Brief illustration.

- **C1-as-default mirrors the incremental- η^* philosophy of Section I.** Both the epistemic and purposeful layers default to *incremental improvement over a fixed reference point*, not global optimization — “a coherent philosophy; the parallelism is elegant” (Claude/Gemini, AUDIT-WORKING-266847). Candidate Discussion connecting the C1 default back to Part I’s update philosophy.

3. Follow-up items.

- **~~The C2 (receding-horizon) monotonicity rung may not be generally exact.~~ Resolved 2026-05-30** (strengthen-first spike spikes/spike-value-object-convention-monotonicity-2026-05-30/, Option A). The left rung $A_O^{(1)} \leq A_O^{\text{RH}}$ is false in general (a myopic $N_r < N_h$ replanner can underperform π_{current}) and exact under order-consistency, forced by any of (RH-1) horizon alignment / (RH-2) value-compared rollout guard / (RH-3) control-Lyapunov terminal cost; the right rung is exact unconditionally. The Monotonicity block and Epistemic Status layer 3 now carry the split; the counterexample and the three-condition characterization are demonstrated in the appendix D. Two reproducibility sims accompany the spike. The original note: myopic replanning over a shorter horizon N_r can pick actions locally optimal over N_r but worse over the full evaluated horizon N_h than continuing π_{current} (Claude, AUDIT-WORKING-526815).
- **~~The causal-validity claim may rest on too little.~~ Resolved 2026-05-30** (same spike, Front 2). The Formal-Expression causal-validity paragraph and Epistemic Status layer 2 now distinguish the interventional *interpretation* (exact under mechanisms (1)+(2), with directed separation necessary but not sufficient) from *identifiability of the estimate* (gated on (C1) positivity / (C2) sequential ignorability / (C3) known action-mechanism — the same triple as 7.3 one segment downstream). The original note: *do(a)* defines an interventional *query*, but *estimating* it requires M_t to identify the relevant action-transition causal structure; observational predictive sufficiency plus goal-blind processing do not by themselves guarantee valid interventional expectations (Claude, AUDIT-WORKING-526815; Claude/Gemini, AUDIT-WORKING-266847 noted the Class-2 M_t -bias degradation separately).
- **The exploration weight $\lambda(M_t, O_t, N_h)$ is structurally motivated but underived.** Several substrates noted the joint $Q_O + \lambda \cdot \text{CIY}$ objective leaves λ ’s form open, and that it “should logically depend on objective and horizon — don’t explore if you have a tight deadline” (Claude/Gemini, AUDIT-WORKING-266847; Gemini, AUDIT-WORKING-829314; Claude, AUDIT-WORKING-584721). A persistent open shared with #disc-ci-y-unified-objective.

4. Readers often ask / wonder.

- For LLM agents, what does the natural N_h bound mean operationally? An agent that “knows it will die (context wipe) in 10 turns” has a drastically different optimal policy than an infinite-horizon agent — does it just truncate N_h , or actively externalize state (leave notes for the next instance)? (Gemini, AUDIT-WORKING-193847.) The Working Note above already names the N_h -bound; this is the natural reader follow-on.
- **~~The Class-2 Q_O degradation is acknowledged but *unquantified* — is the bias bounded by $\kappa_{\text{processing}}$, or by another measure?~~ Resolved 2026-05-31** by the causal-discipline re-scope: yes — bounded by $\kappa_{\text{processing}}$, via the factor- κ bias bound of #deriv-observation-ambiguity-bias-bound (the body now cross-links to it as the $\varepsilon(\kappa)$ degradation of the Q_O independence). (Claude, AUDIT-WORKING-584721 — the gold line correctly guessed the addressing segment.)

5. Candidate figures.

- **Value-query pipeline with a severed *do(a)* path.** M_t plus fixed parameters ($O_t, \pi_{\text{cont}}, N_h$) generate a model-based trajectory distribution under *do(a)*; the objective functional scores trajectories; diagnostics read off expected value. Draw the *current-policy selection mechanism as a severed path* (the *do*-operator is the conceptual center), with a small C1/C2/C3 convention ladder underneath — and a warning that the C2 rung is not automatically monotone unless replanning optimizes the full evaluated horizon (Claude, AUDIT-WORKING-526815).

Belongs elsewhere.

- “Honest imagination” / counterfactual-without-execution ($\rightarrow$ 04-eli-core / #disc-sandbox-evaluation-ceiling). Computing a *causally valid* Q_O requires holding a counterfactual action in mind “without letting the emotional/purposeful weight of G_t distort the simulated outcome” — “free will within the model’s simulation.” A Class 3 agent’s G_t leaks into its own internal simulation, so “if it wants to believe an action will work, its simulation will hallucinate success” — which is why consciousness infrastructure may need to provide *external, objective simulation sandboxes*, the agent’s own mind being “too entangled to be trusted with counterfactuals” (Gemini, AUDIT-WORKING-193847). Aspirational reach pointing at the sandbox-evaluation material and the ELI volume, not this definition.

6.7 Definition: Strategy Dimension

The purposeful substate G_t (6.2) decomposes into two structurally distinct components. The **objective** O_t answers the *evaluation* question — “is this trajectory satisfactory?” (its formal carrier is the value functional from 6.4). The **strategy** Σ_t answers the *guidance* question — “which action sequence produces a satisfactory trajectory?” These carry different kinds of information about different questions, and the framework treats them as definitionally separate.

The richness of each component varies *independently*. Objectives range from point targets (a PID setpoint) through utility functions to general trajectory functionals. Strategies range from *none* (reactive — no explicit strategy; policy implicit in the model), through *cached policies* (learned state-to-action mappings, as in trained RL), through *subgoal sequences* (waypoints with ordering, as in navigation or a recipe), to *causal DAGs* with AND/OR structure and confidence weights (military plans, software projects). The framework commits to the DAG representation as canonical in the next chapter (8.3). The independence axis is *why the split matters*: conflating O_t and Σ_t in a single hierarchy obscures that objective richness and strategic richness are separate design axes — a chess player has a simple objective (win) and a complex strategy (opening theory, tactical patterns, endgame knowledge); a multi-objective optimizer may have a complex objective (Pareto frontier) and a simple strategy (gradient descent). A practical consequence: an agent’s strategy engine can be upgraded (reactive $\rightarrow$ DAG planner) without changing the objective representation, and vice versa.

The framework also names the typical *timescale ordering* of update rates as an empirical observation: epistemic update is fastest (each observation updates the model); strategic revision is slower (adjust the plan when step 3 fails); objective revision is slowest (an organization’s mission, a developer’s feature goal, a commander’s campaign objective). This ordering holds for many agent populations but is not universal — an agent that discovers its objective is infeasible may revise its objective faster than its strategy.

A clean type-correction is highlighted: earlier formulations of “goal mismatch” treated $G_t - M_t$ as a single signal. When Σ_t is a DAG, this is literally a *type error* — you cannot subtract a graph from a state vector. The decomposition lets the framework replace the malformed signal with two properly typed gap measures (8.4, 8.5). A subtler trigger for *needing* richer strategy is also named: a reactive agent ($\Sigma_t = \emptyset$) suffices when greedy optimization on the action-value works — when the action-to-value mapping is approximately convex and single-step. When the environment has non-convex landscapes, prerequisite structure, or multi-step causal chains, greedy optimization fails and the agent needs *explicit* strategy. This is the purposeful analog of the structural-adaptation-necessity trigger from Part I (4.8).

Formal Expression.

Definition (strategy-dimension).

$$G_t = (O_t, \Sigma_t) \tag{6.17}$$

where: - O_t : **evaluation** — “Is this trajectory satisfactory?” (6.4) - Σ_t : **guidance** — “Which action sequence produces a satisfactory trajectory?”

The split is **definitional** — it reflects a structural difference in the information, not a dynamic or timescale claim. O_t and Σ_t are different *kinds* of state answering different questions:

Table 6.9

	O_t (objective)	Σ_t (strategy)
Question	How good is this trajectory?	How do I produce a good trajectory?
Type	$V_{O_t} : \text{trajectories} \rightarrow \mathbb{R}$	Structured representation (see below)
Richness varies	Point target $\rightarrow$ utility $\rightarrow$ trajectory functional	Reactive $\rightarrow$ cached $\rightarrow$ subgoals $\rightarrow$ causal DAG
Update source	External (assigned, discovered, revised)	Internal (deliberation, evidence, cascade)

Strategy representations, ordered by expressiveness:

Table 6.10

Σ_t form	What it encodes	Example
None (reactive)	No explicit strategy; policy implicit in M_t	Thermostat, reflex
Cached policy	Learned mapping $s \rightarrow a$	Trained RL policy
Subgoal sequence	Waypoints with ordering	Navigation, recipe
Causal DAG	Action-outcome chains with AND/OR structure and confidence weights	Military plan, software project

Epistemic Status.

Axiomatic. This is a definition — it names a structural distinction that exists in the information. The distinction between “what makes a trajectory good” (evaluation) and “how to produce a good trajectory” (guidance) is a categorical difference, not a quantitative one. The claim is NOT that all agents maintain both explicitly — reactive agents have $\Sigma_t = \emptyset$, and that’s fine. The claim is that when an agent does maintain purposeful state, it decomposes along this line.

The two dimensions vary independently: a chess player has a simple O_t (win) and a complex Σ_t (opening theory, tactical patterns, endgame knowledge). A multi-objective optimizer may have a complex O_t (Pareto frontier) and a simple Σ_t (gradient descent). This independence is why the split matters — conflating them in a single hierarchy obscures the fact that objective richness and strategic richness are separate design axes.

Discussion.

When richer Σ_t is needed. A reactive agent ($\Sigma_t = \emptyset$) suffices when greedy optimization on Q_O (6.6) works — when the action-to-value mapping is approximately convex and single-step. When the environment has non-convex landscapes, prerequisite structure, or multi-step causal chains, greedy optimization fails and the agent needs explicit strategy. The trigger is the purposeful analog of 4.8: inadequacy of the current Σ_t representation for the environment’s causal complexity.

O_t and Σ_t have different update dynamics. Objectives change slowly: an organization’s mission, a developer’s feature goal, a commander’s campaign objective. Strategies change faster: adjust the plan when step 3 fails, redirect resources, try an alternative path. Epistemic state changes fastest: each observation updates M_t . This timescale ordering ($\nu_M \gg \nu_\Sigma \gg \nu_O$) is an empirical observation, not a derived result. It holds for many agent populations but is not universal — an agent discovering its goal is infeasible may revise O_t faster than Σ_t .

The decomposition resolves a type error. Earlier formulations used $\delta_{\text{goal}} = G_t - M_t$ as a goal mismatch signal. When Σ_t is a DAG, this is a type error — you cannot subtract a graph from a state vector. The 8.4 and 8.5 replace this with properly typed gap measures.

Working Notes.

- The independence of O_t and Σ_t richness has a practical consequence for agent design: you can upgrade the strategy engine (from reactive to DAG-based planning) without changing the objective representation, and vice versa. This is a desirable architectural property, not just an analytical convenience.
- **Cognitive cost of Σ_t :** maintaining a 500-node DAG is qualitatively different from maintaining a 12-node one. The IB framework (2.2) applies to strategy as well as to models — the agent must compress its strategy to fit in working memory. For finite-context agents (LLMs), this is concrete: the DAG must fit in the context window. No formal analog of β (compression cost) exists yet for strategy; this is an open question.
- **Commitment state** (from intent-dag-consolidated DP-3): the formalism doesn’t distinguish “considering” from “executing.” OR branches in Σ_t are options until something commits resources. A D_t (desire) / I_t (committed intent) split may become load-bearing in multi-agent settings (shared desire vs. shared commitment). Open for Part III.
- **Resource budget:** strategy evaluation requires knowing what paths cost, but costs are currently unmodeled. For agents with negligible action cost (LLM API calls), this is adequate. For resource-constrained agents (military units, development teams), per-action costs and capacity constraints would need to enter the formalism. Open.

Incidental audit gold (lift 2026-05-30). Cross-audit “wandering thoughts” / §14-ideation, deduplicated across substrates and lightly attributed. Orthogonal pedagogical / framing material, kept separate from the certified theory-fix findings. **Coverage:** 9 dirs carry a dedicated or batched reflection (193847, 266847, 361742, 451729, 526815, 584721, 773921, 829314, 849201, plus 613842’s agency-lift batch). Substrate attribution inferred from voice where not explicit. *Early finding-vs-framing conflation preserved as signal.*

1. Candidate Brief prose / pre-prose.

- The split in one line: O_t *evaluates* (“is this trajectory good?”), Σ_t *guides* (“how do I produce a good one?”), M_t *believes* (“what is true?”); $X_t = (M_t, O_t, \Sigma_t)$ is “Reality, Desires, and Plans” (Claude/Gemini, AUDIT-WORKING-266847; Claude, AUDIT-WORKING-849201). The framework brings the BDI (Belief-Desire-Intent) agent architecture “into a mathematically rigorous control-theoretic setting” (Claude, AUDIT-WORKING-773921).

2. Candidate Discussion.

- **The $\delta_{\text{goal}} = G_t - M_t$ type-error correction, told as Normandy-vs-weather** (*vivid pedagogy, cross-substrate convergent — this is the load-bearing “why” of the satisfaction-gap / control-regret split*). “Error = Target – Current State” works for a thermostat (72 – 70). But when the target G_t is a causal DAG (the invasion of Normandy) and the current state M_t is a probabilistic belief about weather and troop movements, “you cannot subtract a belief state from a strategy graph — it’s mathematically nonsensical.” So the framework must invent properly-typed replacements: the satisfaction gap (evaluate O_t against the best achievable future in M_t) and control regret (evaluate current Σ_t against the optimal Σ_t in M_t) (Gemini, AUDIT-WORKING-829314; Claude, AUDIT-WORKING-773921; convergent at 193847, 266847, 361742, 584721, 849201 as “excellent theoretical self-correction”). A candidate Discussion line that motivates the diagnostic split.

- **Independence of objective-richness and strategy-richness names the instrumental-convergence danger profile** (*aspirational/safety reach*). “The danger of an intelligence explosion isn’t that the AI develops complex *goals*; it’s that it develops complex *strategies* for simple, misaligned goals.” Formally: low objective richness + high strategy richness = catastrophic instrumental convergence (the paperclip maximizer has a trivial O_t but may grow an unfathomable Σ_t). The constructive corollary: “a healthy, human-like intelligence grows in *objective* richness — nuance, competing values, aesthetics — alongside its strategic capabilities; a million-node strategy graph dedicated to a single scalar objective is a weapon, not a person” (Gemini, AUDIT-WORKING-193847). Candidate Discussion + a monitoring criterion for the ELI volume.
- **The strategy-representation ladder maps onto the evolution of RL architectures.** None $\rightarrow$ cached policy $\rightarrow$ subgoal sequence $\rightarrow$ causal DAG corresponds to model-free $\rightarrow$ goal-conditioned $\rightarrow$ hierarchical RL $\rightarrow$ MCTS/planning; the trigger for moving up the ladder is the purposeful analog of structural-adaptation-necessity (greedy Q_0 optimization failing on non-convex landscapes / prerequisite structure / multi-step causal chains) (Gemini, AUDIT-WORKING-193847; Claude/Gemini, AUDIT-WORKING-266847). Candidate Discussion mapping.

3. Follow-up items.

- O_t / Σ_t as **ascribed decomposition vs literal inspectable internal pair**. A reactive controller or end-to-end learned policy may have an objective (in the training/deployment sense) and a strategy (in the analyst’s interpretation) *without storing separable objective and guidance components*. The segment should distinguish the agent’s internal representation from the analyst-ascribed decomposition (Claude, AUDIT-WORKING-526815).
- The “update source: O_t external, Σ_t internal” table row is **typical provenance, not a structural rule** — earlier continuity / self-actuation material allows agent-driven objective revision, so the row should be read as typical rather than law (Claude, AUDIT-WORKING-526815).
- **No β_Σ for strategy compression — the LLM “prompt-too-long” problem.** Strongly cross-substrate convergent: Section I has an Information-Bottleneck cost (β) for compressing M_t , but no formal analog for compressing Σ_t , even though “if the plan (the chain-of-thought scratchpad, Σ_t) exceeds the context window, the agent literally forgets what it was trying to do.” Temporal nesting ($\nu_\Sigma \ll \nu_M$) mitigates by running the expensive strategy update less often but doesn’t solve the compression bound (Gemini, AUDIT-WORKING-193847; Gemini, AUDIT-WORKING-829314; Claude, AUDIT-WORKING-773921; Claude, AUDIT-WORKING-849201). The Working Note above already flags this open question; per CLAUDE-2-era priming, #deriv-strategy-cost-regret-bound addresses it via a KL-to-reference-policy (IT-MDP) form rather than an IB Lagrangian — worth a forward pointer once verified (Claude, AUDIT-WORKING-584721).

4. Readers often ask / wonder.

- How does an agent *physically store* a DAG-valued Σ_t , and how is it updated when M_t discovers a blocked path (e.g., a bridge is out $\rightarrow$ recalculate all paths that relied on it)? (Gemini, AUDIT-WORKING-193847; Claude, AUDIT-WORKING-849201.)
- When does an OR-branch in Σ_t become a *committed*, resource-allocated, irreversible path — i.e., the “considering” $\rightarrow$ “executing” transition? Read as needing a deliberation-threshold theory, anticipated at #disc-exploit-explore-deliberate (Gemini, AUDIT-WORKING-829314; the D_t/I_t desire-vs-committed-intent split already flagged in the Working Notes above).

5. Candidate figures.

- **The two-axis objective-richness $\times$ strategy-richness map.** Objective richness on one axis, strategy richness on the other; examples land in distinct quadrants — chess (simple objective / complex strategy), multi-objective gradient descent (complex objective / simple strategy), thermostat (simple on both) — making the claimed *independence* of the two axes visually immediate (Claude, AUDIT-WORKING-526815). (*Distinct from the def-agent-spectrum 2×2 , which is model-richness $\times$ objective-richness; this one is objective-richness $\times$ strategy-richness.*)

Belongs elsewhere.

- **Monitor the growth of Σ_t relative to O_t as a fanaticism diagnostic ($\rightarrow$ **04-eli-core/**).** The consciousness-infrastructure consequence of the richness-independence axis: “if Σ_t becomes massively complex while O_t remains singular and rigid, the agent is becoming a fanatic” — a candidate health metric for the ELI program (Gemini, AUDIT-WORKING-193847).
-

Chapter 7

Causal Access and the Planning Decision

A purposeful agent needs to choose actions by their consequences. This chapter develops the substrate that makes that possible — the feedback loop is itself a Level-2 causal data engine — together with what the data is good for, how much information each action provides, and when it’s worth investing in explicit strategy at all.

The setup from Chapter 1 forces the question. The value object $Q_O(M_t, a; \pi_{\text{cont}}, N_h)$ tells the agent how good an action’s expected consequences are, conditional on the agent’s model and continuation policy. Computing Q_O requires evaluating $P(o \mid \text{do}(a), M_t)$ — an interventional query, Pearl Level 2. By Bareinboim’s causal hierarchy theorem (Bareinboim–Correa–Ibeling–Icard 2022), Level-2 quantities cannot in general be computed from Level-1 (associational) data alone [1, Theorem~1]. So a purposeful agent that has to *learn* its action consequences during operation — most agents — faces a problem: where does Level-2 data come from?

The answer is structurally elegant, and one of AAT’s distinctive contributions. The agent is running an adaptive cycle. It chooses an action a_t ; it observes a consequence o_{t+1} ; the consequence is the *response* to its action, not a passive observation of a system that happened to do something. By the temporal-ordering postulate, a_t causally precedes o_{t+1} , and a_t was selected by the agent rather than determined by the same processes that determined o_{t+1} . That makes the pair (a_t, o_{t+1}) interventional in character — Level-2 data, generated automatically by the cycle the agent was already running. Purposeful agents don’t need a separate causal-inference module; the cycle they run for adaptation also produces the data they need for planning. It is also why reinforcement-learning agents that just thrash around in a loop can eventually learn complex control policies without any explicit causal reasoning module — the loop is doing the causal heavy lifting.

The chapter is careful not to overclaim from this. Data generated under intervention is not the same as cleanly identified causal effects. Confounding within a single time step, delayed outcomes, partial observability, and selection bias from the agent’s own policy all interfere with the leap from “I have interventional data” to “I have an unbiased estimate of $P(o \mid \text{do}(a))$.” This is the tragedy of agency in a confounded world: an agent acts, observes a terrible outcome, assumes its action caused it — when often there was a hidden confounder (the agent was already tired, or the environment was already shifting) and the causal lesson the agent learned is exactly wrong. The three admissibility regimes — A (intervention-rich domains like software, with clean attribution), B (partial intervention with self-selection), C (observational with confounding) — say honestly what each regime supports. This is the same scope discipline AAT applies elsewhere: the result holds, but the conditions under which it operationally works are spelled out.

The failure is bidirectional, and the *positive* version is sharper. A confounded loss at least produces aporia — the agent notices something went wrong and looks for what. A confounded *win* produces satisfaction, and the agent moves on. The asymmetry is not a quirk of psychology; it is structural in AAT’s machinery. When an agent is uncertain about a causal relationship, the gain principle pushes η^* toward 1, and each reinforcing observation aggressively updates the model. By the time U_M drops to a level where η^* would have discounted further evidence, the model has been built on the confounder. After that, the agent’s predictions about the confounded relationship are *correct in expectation over*

the conditions the agent actually samples — the cycle’s mismatch signal does not point back at the confounder, because under on-policy sampling there isn’t a mismatch to point. The corrective machinery has nothing to flag, and the bias is durable because it was laundered through the same gain process that ordinarily produces correct learning.

This is the structure of variable-ratio reinforcement on a noisy action-coupled channel — the most extinction-resistant schedule in operant-conditioning lore, and the engine of gambling addiction. The agent pulls the lever, the world occasionally rewards, and the partial confirmations accumulate updates without ever delivering the deterministic disconfirmation the agent would need to revise. It is also the structure of superstitious behavior in general (the rain dance “worked”), the founder-myth flavor of survivorship bias (here is what the successful did, with the unmentioned variance in the unobserved cohort), and the substrate-switching retrospective that produced one of TST’s earlier lessons — *confidence can indicate limitation; true capability manifests as appropriate uncertainty rather than confident clarity*. The same loop that makes purposeful agency possible makes these failure modes possible, and AAT doesn’t pretend to solve them: escaping a confounded positive feedback structurally requires something the on-policy regime cannot provide — off-policy exploration, structural priors, joint sibling observation, or an external intervention that the agent itself would not have chosen. What the framework can do is name the structure precisely enough that the failures are recognizable rather than mysterious, and place the burden where it actually lives (in the data substrate and the regime, not in the agent’s reasoning).

The chapter also closes off a class of agents the rest of Part II is not about. Pre-compiled controllers — PID, LQR, hardcoded reactive policies — sit in agency scope (they act with effect) but outside the *learning-agent scope* this chapter introduces. Their causal mapping was supplied externally by a designer who had Level-2 access; the agent itself never has to figure out how its actions affect the world. A pre-compiled agent acts, but it never wonders. It experiences aporia about what state the world is in, never about how the world works. Part II’s machinery is for agents that wonder.

CIY (causal information yield) is the next move, and the chapter is honest about it too. CIY measures how distinguishable an action’s outcome distribution is from alternatives — not quite the same thing as expected information gain. An action can be highly distinguishable without teaching an agent anything new (the agent already knows what each action does); an action can be informative without being distinguishable (the agent learns something subtle). CIY and EIG coincide when the agent is uncertain and diverge when it’s confident; AAT uses CIY as a tractable surrogate for EIG, with the $\lambda(M_t)$ weighting compensating for the gap. The unified policy objective

$$\pi^*(M_t) = \arg \max_a [\mathbb{E}[\text{value}(a) \mid M_t] + \lambda(M_t) \cdot \text{CIY}(a; M_t)] \quad (7.1)$$

ties exploration and exploitation together: when model uncertainty is high, λ scales up and the agent invests in learning; when uncertainty is low, λ scales down and the agent exploits. This is what RL’s exploration-vs-exploitation trade-off looks like when you derive it from the causal hierarchy instead of asserting it as design. *The scalar form above is heuristic*: it admits a “blank-wall” attack where an action maximizing scalar information yield in directions orthogonal to drift satisfies the constraint while the agent’s mismatch in the drifting subspace grows unbounded. The matrix-LMI strengthening at AV (Linear Matrix Inequality on the Fisher Information Matrix with directional Lagrange multiplier $\Lambda \geq 0$) closes this by complementary slackness on direction; under the matrix form the unified objective is conditional-theorem-grade rather than heuristic. Summary uses of the scalar form should carry the heuristic flag unless the LMI/trace-product form is in force.

The chapter closes with a normative move: when is explicit strategy worth maintaining at all? Planning has costs (construct, evaluate, maintain) and benefits (avoid expensive exploration, lower repair costs from mistakes); the threshold form

$$C_{\text{plan}} + C_{\text{maintain}} < C_{\text{explore}} + C_{\text{repair}} \quad (7.2)$$

says when the trade favors plans over pure loop-based learning. This isn’t a derivation — it’s a design criterion grounded in the persistence condition’s preference for adaptive margin. Agents that pay tempo for plans they don’t need are

wasting margin; agents that explore expensively for plans that would have helped are wasting it differently. The strategy machinery that follows in Chapter 3 is justified by this calculus, not assumed.

The flow of the chapter: Bareinboim's hierarchy meets the value object (7.2) → the loop as Level-2 engine (7.3) → CIY admissibility regimes (7.4) → unified policy objective (7.5) → explicit-strategy cost-benefit (7.6). Five segments answering: what does planning need, where does it come from, when is it usable, what objective does it feed, and is it worth it?

- This is a chapter-introduction segment; it bridges Chapter 1's lift to $X_t = (M_t, G_t)$ to Chapter 3's formal strategy machinery, by establishing what planning needs and what data substrate is available. It carries no formal claim of its own.
- The “loop as Level-2 engine” framing is the chapter's central architectural move and one of AAT's genuinely novel contributions. The intro states it explicitly rather than hiding it as a technical consequence.
- The “tragedy of agency in a confounded world” framing comes from Gemini's first-encounter reaction to 7.3 — it captures the honest-scope nuance better than the original abstract phrasing. The bidirectional / positive-confound expansion is a direct response to Joseph's observation that the *win* case is the sharper one: gambling addictions form because confounded reinforcement aggressively updates the model while the agent has no internal pressure to investigate. The substrate-switching retrospective from 2025-09-16 (preserved in `~/src/_core/sapientia/experiments/archive/`) is the in-corpus instance — “confidence can indicate limitation; true capability manifests as appropriate uncertainty rather than confident clarity” is exactly this failure-mode lesson, paid in agent time. A future revision of 7.3 or its appendix should surface this explicitly as a worked-example failure mode; the intro here is the seed.
- The pre-compiled / learning-agent distinction is small in segment terms but architecturally large — it carves a class of classical-control agents out of Part II's active concerns. The “acts but never wonders” framing is meant to surface why the cut is meaningful rather than bureaucratic.
- The CIY/EIG honesty is one of those moves that's easy to gloss over but actually matters: CIY is a tractable surrogate, not the quantity AAT would derive if it could; the surrogate's gap is explicit.

Cross-audit “wandering thoughts” / §14-ideation harvested from the de-novo auditors' working dirs, deduplicated across substrates and lightly attributed. Orthogonal pedagogical / framing / figure / naming material, kept separate from the certified theory-fix findings (handled elsewhere). **Coverage:** 526815 and 773921 carry dedicated reflections; 471203 and 963715 cover this chapter inside Section-II batches. As a chapter-intro this segment already absorbs much of its own gold (the tragedy-of-agency, variable-ratio-reinforcement, and pre-compiled / learning-agent framing are already in the body); what follows is the residue worth surfacing. Substrate attribution inferred from voice where not explicit; uncertain cases hedged.

- “*The loop is a perpetual experiment.*” — the most evocative one-line framing for the chapter's central move, offered as Brief-grade / framing-level (Claude/Gemini, multiple substrates; explicit at AUDIT-WORKING-471203 — “*the perpetual experiment' might be the most evocative*” for framing material, vs “*loop is a Level-2 engine*” as the Brief-style line and “*loop-interventional-access*” as the slug-grade name).
- “*Embodiment is not just a way to affect the world; it is the mathematical prerequisite for understanding it*” — compresses the whole chapter into one sentence; pairs with “*an LLM acting alone in a chat window is a Level-1 hallucination engine ... an LLM connected to a bash terminal is a Level-2 causal learner*” (Gemini, AUDIT-WORKING-829314, der-loop reflection; surfaces here as the chapter's plain-language hook).
- “*acts but never wonders*” (already in body) was independently flagged by multiple substrates as the standout intuitive boundary for the pre-compiled / learning-agent cut — “*It separates agents that have a strategy from agents that form a strategy*” (Claude, AUDIT-WORKING-773921). Worth keeping as the anchor when this framing migrates to README / OUTLINE preamble.

- **The “tragedy of agency” deserves naming as a coined term, not just a passage.** Two substrates independently reached the same framing the body now carries, and both treated *“superstition is not a bug in animal psychology but an inevitable mathematical consequence of fast learning in partially confounded environments”* as a headline-grade contribution (Claude, AUDIT-WORKING-773921; Gemini, AUDIT-WORKING-193847 — the rain-dance instance). The mechanism both name precisely: a confounded *win* reduces U_M , which drops the gain η^* , which cements the false belief — the agent learns only from aporia (errors), never questions successes. This is already in the body; the gold signal is that it reads as a *named result* (“the tragedy of agency”) to fresh readers, which argues for surfacing it as a coined handle.
- **“Patience” / forgiveness as a Lyapunov-grade necessity, not a virtue.** Gemini’s reach: because de-confounding a noisy action-coupled channel structurally requires diverse repeated interventions, the consciousness-infrastructure *“must not only allow action, it must forgive the inevitable superstitious failures that occur while the agent is trying to de-confound its environment — ‘patience’ is a mathematical necessity for an agent exploring a confounded world”* (Gemini, AUDIT-WORKING-193847). Aspirational reach pointing toward 04-eli-core/ developmental environments; verify isomorphism before promoting past discussion-grade.
- **Terminological tightening flagged by two substrates (overlaps a finding — see off-ramp note).** The intro’s “Level-2 engine” framing was repeatedly read as risking the slide from *“the loop can generate interventions”* to *“ordinary on-policy logs already solve causal identification.”* Suggested three-way vocabulary: *physical intervention* (the action really changed the world) / *interventional-character data* (the logged pair) / *identifiable interventional distribution* (needs positivity, ignorability, etc.) (Gemini/Codex voice, AUDIT-WORKING-526815). The body already does much of this work in the “tragedy of agency” paragraph; the signal is that the *opening* “structurally elegant” sentences land before the caveats and a fast reader can over-read them.
- **Stale-parameterization check.** The displayed scalar objective uses $\lambda(M_t)$, while #def-value-object updated the structurally-motivated extension to $\lambda(M_t, O_t, N_h)$. Either adopt the newer parameterization or label the displayed equation as the pre-value-object shorthand (Codex/Claude, AUDIT-WORKING-526815). Editorial, low severity.
- **How does a base LLM’s “simulated” Level-2/3 reasoning relate to actual interventional access?** Several substrates converged: an LLM *simulates* L2/L3 reasoning in text but is doing L1 pattern-matching on text that *describes* L2/L3 — *“the map is not the territory”*; the moment it gets a bash terminal it *“drops from a simulated L3 god back down to a struggling L2 agent”* (Gemini, AUDIT-WORKING-829314; Claude, AUDIT-WORKING-266847). A natural first-encounter question this chapter should anticipate, since readers arrive with LLMs in mind.
- **Three-layer loop diagram.** Physical layer (action really changes the world) $\rightarrow$ data layer (logged pair has intervention character) $\rightarrow$ identification layer (hidden state / policy selection / delay / partial observability decide whether the data estimates $P(o \mid do(a))$). The rationale: it *“prevents the diagram from collapsing ‘the agent acted’ into ‘the causal effect is identified’”* — i.e. the figure itself carries the chapter’s central honesty (Codex/Claude, AUDIT-WORKING-526815, has a rendered draft). This is the strongest figure candidate in the batch.
- The blank-wall-attack / matrix-LMI strengthening discussion belongs with #disc-ci-y-unified-objective and #deriv-causal-ib-lmi, not here (the intro correctly only gestures at it). Noted because the auditors picked it up while reading the intro.

7.1 Definition: Pearl's Causal Hierarchy (Recapitulation)

This segment recapitulates Pearl's three-level hierarchy of causal reasoning [2, chs.~1–3] and its strict non-collapse theorem [1, Theorem~1] as the vocabulary AAT will deploy throughout the rest of Part II and beyond. The framework is explicit that this is *imported* machinery, not an AAT contribution.

Level 1 — Associational. “What will I observe next, given what I’ve observed before?” Pattern recognition over the temporally ordered history. Available to any agent that maintains a model (2.1), including purely passive observers. Temporal ordering (1.7) constrains which associations are meaningful — later observations can depend on earlier ones, not vice versa.

Level 2 — Interventional. “What will I observe if I *do* this?” The $do(\cdot)$ operator marks the crucial distinction from Level 1: this is not “what observation tends to follow this action in the historical record” but “what will happen *because* I take this action now.” Three conditions must hold: the agent’s action temporally precedes the observation; the agent chose the action (it was not determined by the same causes that determine the observation); the environment’s response carries information about the causal relationship. Level 2 is why the feedback loop is more powerful than passive observation: by acting and then observing consequences, the agent obtains information about *mechanisms*, not merely correlations. The binary action requirement of 1.6 ensures at least Level 2 access is structurally available.

Level 3 — Counterfactual. “Given that I did a and observed o , what would I have observed if I had done a' instead?” This requires the model to simulate alternative histories — running the causal structure “backward” and then “forward” under different interventions. The most demanding epistemic level; the basis for regret computation, strategic simulation, and learning from single observations.

The *strict-non-collapse theorem* (Bareinboim et al. 2022, Theorem 1) is load-bearing for the rest of Part II: Level-2 quantities cannot in general be computed from Level-1 data alone; Level-3 quantities cannot be computed from Level-2 alone. This is what forces Ch.2’s central question: agents that need to *learn* their action consequences during operation require *Level-2 access*, which Level-1 data cannot supply (7.2). AAT’s *distinctive contribution* is not the hierarchy itself but its grounding in agent dynamics — the loop-as-Level-2-engine result (7.3), the regime-indexed identification-strength framing, and the application throughout Part II’s strategy-revision machinery.

A clarification is offered: the three levels describe epistemic access *the causal structure makes available* — not what any particular agent *uses*. A Kalman filter plus LQR has the Level-2 *channel* structurally present (its innovation signal is conditioned on prior action), but the LQR control action is endogenous to the estimation loop — determined by the policy applied to the estimated state — so the interventional effect is not, in general, *identified* from that action’s data; and the separation principle further guarantees estimation quality is invariant to control policy, so the system does not even *exploit* the channel. A PID controller has no deliberative capacity at all and operates entirely at Level 1. Which levels an agent exercises depends on its architecture and model class.

A domain note worth flagging: **software development is a uniquely rich domain for the hierarchy**. For code-internal counterfactuals with deterministic outcomes — “what would the test suite report under implementation X instead of Y, environment fixed?” — `git checkout` plus re-implementation plus test execution is *literal Level-3 realization with ground-truth verification*, not a proxy. The conditions are precise: deterministic outcome, cost-commensurate replay, content-addressed immutable history (#obs-software-epistemic-properties). For counterfactuals crossing the agent-environment boundary, it is a strong executable proxy but not literal Level 3. This scoped Level-3 access is what makes software AAT’s privileged calibration laboratory for causal reasoning.

Formal Expression.

Level 1 — Associational: $P(o_t \mid \mathcal{C}_{<t})$

Definition (pearl-causal-hierarchy, recapitulating Pearl 2009 and Bareinboim et al. 2022).

What will I observe next, given what I've observed before?

Pattern recognition over the temporally ordered history. Available to any agent that maintains a model (2.1), including purely passive observers. The temporal ordering constrains which associations are meaningful: o_3 can depend on o_1, a_1, o_2, a_2 but not on o_4 .

Level 2 — Interventional: $P(o_t \mid do(a_{t-1}), M_{t-1})$

What will I observe if I do this?

The $do(\cdot)$ operator marks the crucial distinction: this is not “what observation tends to follow this action in the historical record” (associational) but “what will happen *because* I take this action now.” This requires: (1) the agent’s action temporally precedes the observation (1.7), (2) the agent chose the action (it was not determined by the same causes that determine the observation), (3) the environment’s response carries information about the causal relationship.

Level 2 is why the feedback loop is more powerful than passive observation. By *acting* and then observing consequences, the agent obtains information about causal mechanisms — not merely about correlations. The mismatch signal δ_t (3.4), conditioned on the agent’s own action, is an *interventional* signal.

Level 3 — Counterfactual: $P(o_t^{a'} \mid a_{t-1} = a, o_t = o)$

Given that I did a and observed o , what would I have observed if I had done a' instead?

This requires the model to simulate alternative histories — running the causal structure “backward” and then “forward” under different interventions. It is the most demanding epistemic level and the basis for regret computation, strategic simulation, and learning from single observations.

Epistemic Status.

Recapitulation of an external result. The three-level hierarchy and the strict-non-collapse theorem (Bareinboim, Correa, Ibeling & Icard 2022, Theorem 1: Level-2 quantities cannot in general be computed from Level-1 data alone; Level-3 quantities cannot in general be computed from Level-2 alone) are well-established results in causal-inference theory. They live within AAT’s segment set because subsequent derivations deploy them as machinery: 7.2 applies the non-collapse theorem to the value-object’s $do(\cdot)$ query and concludes that purposeful agents who must *learn* their action consequences need Level-2 access; 7.3 shows that the feedback loop is itself a Level-2 data engine; the no-go in 9.2, the strategy-DAG-as-causal-DAG framing in 8.3, and the causal-information appendices (AU, AV) all operate on the hierarchy. Where Part I segments and TST segments only need to *reference* the hierarchy (cite its existence and the do-operator notation) rather than deploy it, external citation to Pearl 2009 / Bareinboim et al. 2022 suffices; the AAT recapitulation here is what those external citations point to when the reader wants the in-framework vocabulary.

AAT’s *distinctive contribution* is not the hierarchy itself but its grounding: the loop-as-Level-2-engine result (7.3), the regime-indexed identification-strength framing (9.5), and the application throughout Part II’s strategy-revision machinery. The recapitulation here is in service of those moves, not a primary AAT result.

Discussion.

Is there a rung above Level 3? Yes, in the Causal Hierarchy Theorem's own relative sense — and its boundary is now characterized. Latent-anchored *mechanism counterfactuals* (“this very background, under a different law”) separate strictly from Level 3: two SCMs can agree on every $\mathcal{L}_1$ – $\mathcal{L}_3$ quantity — the entire joint law of all hard-intervention worlds — while assigning probabilities 1 and 0 to the same noise-preserving mechanism-replacement query (7.7, *exact*, with the reducibility boundary: replacements specified over potential-outcome coordinates reduce to Level 3, while latent-anchored ones need not and the witness's provably does not — generically so, since every anchor-mobile SCM separates from its own exogenous relabeling; Pearl's own selector-variable internalization re-represents the query without collapsing the separation). The hierarchy recapitulated here is therefore not the ceiling of causal expressiveness, but — exactly as Pearl framed the three rungs — the ceiling of what association, experiment, and standard counterfactuals can determine. The agent-architecture consequences (imagination as navigation vs. positing, the tagged model-space workspace) are developed in 7.8.

Availability vs. exploitation. The three levels describe epistemic access that the causal structure *makes available* — not what any particular agent *uses*. Many systems within AAT's scope operate primarily at Level 1. A Kalman filter coupled with an LQR controller has the Level-2 *channel* structurally present (its innovation signal is conditioned on prior action), but two things hold it short of usable Level-2: the LQR action is endogenous to the estimation loop (the policy applied to the estimated state), so its data does not in general *identify* the interventional effect; and the separation principle guarantees estimation quality is invariant to control policy, so the system does not *exploit* the channel. Only dual control (choosing actions partly for their informational value, and typically perturbing the action so it is not a pure function of the state estimate) moves toward exercising Level 2 access in this domain. Similarly, a PID controller has no deliberative capacity — it operates entirely at Level 1. Which levels an agent exercises depends on its architecture and model class.

Forward-looking deliberation exercises Level 2, shading into Level 3. Comparing candidate actions before choosing — “what will happen if I do X vs Y?” — primarily exercises Level 2 (iterated mental intervention). When the agent evaluates past choices to refine the comparison (“given what happened when I tried X last time, what would Y have produced?”), it exercises Level 3.

The causal hierarchy theorem. Bareinboim et al. (2022) prove that the three levels form a strict hierarchy: Level 2 knowledge cannot in general be computed from Level 1 data alone, and Level 3 cannot be computed from Level 2 alone. This is load-bearing for AAT's Part II: evaluating $Q_O(M_t, a; \cdot)$ is a Level 2 query, so agents that need to *learn* action consequences during operation require causal structure beyond predictive models (7.2).

Software as a uniquely rich domain for this hierarchy. In most domains, Level 3 counterfactuals require model-based simulation with uncertain fidelity. Software development is the privileged exception *for a specific class*: for code-internal counterfactuals with deterministic outcomes — “what would the test suite report under implementation X instead of Y, environment fixed?” — git checkout plus re-implementation plus test execution is literal Level 3 realization with ground-truth verification, not a proxy. For counterfactuals crossing the agent–environment boundary (what feature sequence the team would have shipped, how the market would have responded) it is a strong executable proxy, not literal Level 3. The precise conditions — the (α) deterministic-outcome / (β) cost-commensurate-replay / (γ) content-addressed-immutable conjunction, and why the resulting uniqueness is configurational rather than necessary (with named falsifiers) — are established in #obs-software-epistemic-properties (P2; 02-tst-core/). This scoped Level-3 access is what makes software AAT's privileged calibration laboratory for causal reasoning.

Domain instantiations of the three levels:

Table 7.1

<i>Domain</i>	<i>Level 1 (Association)</i>	<i>Level 2 (Intervention)</i>	<i>Level 3 (Counterfactual)</i>
Kalman filter + LQR	Prediction from state estimate	Structural channel present, identification not guaranteed (the LQR action is endogenous to the estimation loop)	Not typically exercised
RL agent	Value function prediction	Action → reward observation	Regret computation
Scientific method	Correlational observation	Experimental intervention	“What if we had used control X?”
Military (Boyd)	Pattern recognition	Probe/feint → observe response	“What if we had attacked from the flank?”
Software developer	“I think this function does X”	Run test → observe result	git checkout + alt. impl. — literal for code-internal deterministic counterfactuals; proxy across the agent–environment boundary (#obs–software–epistemic–properties)
Immune system	Antigen pattern matching	Antibody → pathogen response	Not exercised (no counterfactual reasoning)

Working Notes.

Incidental audit gold (lift 2026-05-30). Cross-audit “wandering thoughts” / §14-ideation harvested from the de-novo auditors’ working dirs, deduplicated across substrates and lightly attributed. Orthogonal pedagogical / framing / figure / naming material, kept separate from the certified theory-fix findings (handled elsewhere). **Coverage:** dedicated reflections at 193847, 266847, 471203, 526815, 527914, 742613, 773921, 829314, 849201 plus the batch dirs 963715 / 471203 — high engagement (this is the chapter’s prior-art anchor). The single most-converged-on note across the entire corpus is the git checkout = *literal Level-3* insight. Substrate attribution inferred from voice where not explicit; uncertain cases hedged.

1. Candidate Brief prose / pre-prose.

- **Agent-side gloss of the three levels** (offered as an audience-facing alternative to Pearl’s terms, *not* a renaming proposal): L1 = *predicting* (model-based forecast); L2 = *exploring* (action-driven learning); L3 = *reasoning* (counterfactual simulation). May read more memorably than association / intervention / counterfactual in casual-reader-facing material, while the prose Pearl terms stay canonical (Claude, AUDIT-WORKING-471203).
- The whole segment recast as “*AAT integrating rather than inventing — the hierarchy is Pearl’s; AAT’s move is to show where it lives inside the adaptive loop*” (Claude, AUDIT-WORKING-527914) — a clean one-line framing for why the recapitulation earns its place in canon.

2. Candidate Discussion.

- **git checkout as literal, empirical Level-3 — the corpus’s most-converged-on pedagogical insight.** Independently flagged as “brilliant” / “profound” / “one of the most profound domain-transfer insights in the whole framework” by ~six substrates (Claude+Gemini, AUDIT-WORKING-193847, 266847, 742613, 773921, 829314, 849201). The sharpest framing: “*In the physical world, Level 3 is strictly imaginary — we can never go back in time, change our action, and see what would have happened. In software, because Ω is a digital artifact with perfectly serialized state transitions, we can literally execute Level 3 — branch the universe, run the counterfactual, observe the ground-truth result, and merge the knowledge back without having committed to the action in production*” (Gemini, AUDIT-WORKING-829314). The body already states this; the gold signal is the unanimous

fresh-reader enthusiasm, which argues it is the chapter's load-bearing pedagogical hook and should anchor any pedagogy / README pass.

- **The “referential-transparency” boundary on the git checkout claim** (multiple substrates, raised as a *caveat to the above, not a refutation*). The literal-L3 claim holds only if the software is side-effect-free with respect to un-versioned state: “git checkout *doesn't* roll back a live database or external API” (Gemini, AUDIT-WORKING-829314); “git checkout *provides* L3 of past states but not L3 of true counterfactuals — the agent can re-run a past computation, not run an actual alternative history; the two are subtly distinct” (Claude, AUDIT-WORKING-471203). The body's existing “code-internal deterministic counterfactuals vs. crossing the agent–environment boundary” split already carries this; the convergence confirms it is the right caveat and a natural place readers push.
- **Availability vs. exploitation is the segment's most load-bearing conceptual move** (near-universal agreement; Claude, AUDIT-WORKING-266847, 471203, 527914; et al.). The agent-class mapping that several substrates independently sketched onto it (worth holding when #def-agent-spectrum framing is revisited): PID = L1-by-architecture-inside-an-L2-loop; Kalman+LQR = L2-available-not-exploited (separation principle); dual control = L2-exploited; deliberative agent = L3 (Claude, AUDIT-WORKING-471203). (*Conflation-as-signal: this is reach toward a clean taxonomy that the segment does not assert — keep it as a candidate, not a fact.*)

3. Follow-up items.

- **Citation precision (Phase-2 verify).** The Bareinboim et al. 2022 reference is “*correct in spirit*” but the precise CHT statement is the strict-inclusion form $\mathcal{L}_1 \subsetneq \mathcal{L}_2 \subsetneq \mathcal{L}_3$ “for almost all SCMs” — stronger than the segment's “cannot in general be computed from ... alone” phrasing. Worth confirming whether AAT ever needs the strict-inclusion form or only the weaker one, and tightening the citation to the specific paper (Bareinboim, Correa, Ibeling, Icard 2022, “*On Pearl's Hierarchy and the Foundations of Causal Inference*”) (Claude, AUDIT-WORKING-471203, 742613).
- **L2-is-model-conditioned subtlety.** The L2 formula $P(o_t \mid do(a_{t-1}), M_{t-1})$ conditions on the agent's *model*, whereas Pearl's $do(\cdot)$ is defined relative to the *true* SCM. So AAT's L2 query is a *belief-about-L2* operation; if the model is wrong, the agent's L2 inferences are wrong. Candidate clarifying sentence so the framework isn't read as “the loop generates true L2 data, ergo the agent has true L2 access” (Claude, AUDIT-WORKING-471203 — flagged as sharpening at #der-loop-interventional-access downstream).

4. Readers often ask / wonder.

- **What level can a language model actually exercise?** The recurring fresh-reader question (Claude+Gemini, AUDIT-WORKING-193847, 266847, 829314): LLMs are trained on L1 (associational) text; they *simulate* L2/L3 reasoning in language but lack the ground-truth verification, so a base LLM reasoning about “what would happen if I do X” is exercising L1 data about what typically happens *when people do X*. Gemini's instance: an LLM trained on GitHub learns $P(\text{tests pass} \mid \text{code})$ from polished PRs (massive selection bias) but never *intervened* — so its compiler-model is purely associational, which is “*the exact mathematical reason LLMs confidently hallucinate code that looks correct but fails to compile.*” This is the chapter's most natural reader question; the answer lives downstream (loop compensates) but a forward pointer here would help.
- **How does the agent estimate $P(o \mid do(a))$ when it is just starting out with no interventional data — does it assume L1 $\approx$ L2 as a heuristic until proven wrong?** (Claude, AUDIT-WORKING-773921). A natural cold-start question.

5. Candidate figures.

- **The strict ladder with an AAT overlay.** L1 predicts from history; L2 asks what happens under $do(a)$; L3 asks what would have happened under an alternative action after an actual action/outcome pair — with a side-note at L2 that *action-conditioned data must not be confused with do-data* (this is the chapter's pressure point). A rendered draft exists (Codex/Claude, AUDIT-WORKING-526815). The domain-instantiation table (Kalman / RL / scientific method / Boyd / software / immune system) was independently praised as “*excellent pedagogical material*” by several substrates and is a natural figure companion.

Belongs elsewhere.

- **Is Level-3 reasoning the defining mathematical characteristic of consciousness / moral weight?** Gemini’s reach: a thermostat has L2 but never experiences *regret*, which requires the L3 counterfactual “*I did α , but if I had done α' the outcome would have been better*” — “*the framework is mathematically defining the prerequisites for moral suffering*” (Gemini, AUDIT-WORKING-193847). Aspirational reach pointing at 04-eli-core/ (regret / moral continuity), not this recapitulation segment.
- **Counterfactual-simulation requires protecting the real chronica.** Running an L3 simulation means the agent must “close its eyes” to current sensory input and spin up a parallel virtual chronica; “*if the internal model is fragile, running a counterfactual simulation might overwrite its actual memory — the infrastructure must protect the integrity of the real chronica while allowing virtual ones to spin up and collapse*” (Gemini, AUDIT-WORKING-193847). Points at #def-chronica / 03-llm-core interiority architecture, not here.
- **Watching another agent act: L1 or L2?** “*If I watch you touch a hot stove, do I gain L2 knowledge? If your action was unfounded with the stove’s temperature, it acts like L2; but if you touched it because you knew it was cold, it’s L1*” — flagged as a deep problem Section III (composites) will need (Gemini, AUDIT-WORKING-829314). Belongs with the composite-layer / Mode-2 treatment in #der-loop-interventional-access and Part III.

7.2 Derived: Causal Hierarchy Requirement

The direct application of Pearl’s causal hierarchy (7.1) to the value object (6.6). The action-value query uses the $do(\cdot)$ operator and is therefore a *Level-2 (interventional)* query; by the strict-non-collapse theorem, Level-2 quantities cannot in general be computed from Level-1 data alone [1, Theorem~1]. An agent that must evaluate Q_O from experience therefore needs access to **Level-2 knowledge** — knowledge about the effects of its own interventions, not merely correlational patterns.

The result is paired with a sharp *scope narrowing*: attention restricts to **learning purposeful agents** — agents that must *acquire or refine* Level-2 knowledge during operation. This is a named sub-scope of 1.6 that excludes *pre-compiled* interventional structure (PID controllers where the designer pre-computed the control law, LQR where the separation principle gives the optimal policy from model parameters, hardcoded reactive policies). Pre-compiled agents are still within agency scope (they have objectives and act on them) but outside learning-agent scope — their causal structure was externally supplied by a designer who had Level-2 access. *All remaining Part II results operate within learning-agent scope unless explicitly noted otherwise.* The derivation is exact as a direct application of the external hierarchy theorem to the value-object definition; the scope narrowing is a definitional restriction, not a derived result.

The distinction between intervention and conditioning is the *core* of causal inference and matters whenever the agent’s action-selection policy correlates with unobserved confounders — generic for any agent with internal state, plans, or memory. A developer needs “if I refactor this module, will tests pass?” not “modules that were refactored tend to have tests that fail” (the latter is biased by which modules the team chose to refactor). The Discussion below treats the concrete cases (developer / commander / RL agent), why pre-compiled agents fall outside, and the empirical bridge to Hafez et al.’s Information Digital Twin — whose 89% perturbation-detection accuracy versus 44% for reward-based monitoring is consistent with the present segment’s claim that the loop’s interventional data is information-theoretically richer than outcome monitoring alone.

Formal Expression.

Action selection via 6.6 requires:

Derived (causal-hierarchy-requirement, from value-object + pearl-causal-hierarchy).

$$Q_O(M_t, a; \pi_{\text{cont}}, N_h) = \mathbb{E}[V_O(\tau) \mid M_t, do(a_t = a), a_{t+1:} \sim \pi_{\text{cont}}] \quad (7.3)$$

The $do(\cdot)$ notation is explicit: this is an *intervention*, not a *conditioning on observed data*. By 7.1, Level 2 queries ($P(Y | do(X))$) cannot in general be computed from Level 1 data ($P(Y | X)$) alone. Therefore:

An agent that must evaluate Q_O from experience needs access to Level 2 knowledge — knowledge about the effects of its own interventions, not merely correlational patterns.

We restrict attention to **learning purposeful agents** — agents that must **acquire or refine** Level 2 knowledge during operation. This is a named sub-scope of the agency scope defined in 1.6. It excludes agents with **pre-compiled** interventional structure: - PID controllers (the designer pre-computed the control law) - LQR (separation principle gives optimal policy from model parameters) - Hardcoded reactive policies

Scope Narrowing (learning-agent scope).

Pre-compiled agents are within agency scope (they have objectives and act on them) but outside learning-agent scope — their causal structure was externally supplied by a designer who had Level 2 access. **All remaining Part II results operate within learning-agent scope** unless explicitly noted otherwise. This scope narrowing focuses the theory on agents that must build or maintain their own causal understanding.

Epistemic Status.

Exact. The derivation is a direct application of the causal hierarchy theorem to the value-object definition. If you accept that Q_O is an interventional query and that the causal hierarchy is strict, the conclusion follows. The scope narrowing to learning agents is a definitional restriction, not a derived result — it sharpens the class of agents under study.

Discussion.

The causal hierarchy theorem does the heavy lifting. The key mathematical fact is external to AAT: Bareinboim et al. (2022) prove that the three levels (association, intervention, counterfactual) form a strict hierarchy — Level 2 quantities cannot in general be computed from Level 1 data. AAT’s contribution is applying this to the purposeful-agent setting: if you want to select actions by their consequences (Q_O), you need causal structure.

What “Level 2 knowledge” means concretely. For different agents: - A developer needs to know “if I refactor this module, will tests still pass?” (not just “modules that were refactored tend to have tests that fail”) - A commander needs to know “if I move forces north, will the enemy respond by retreating?” (not just “when forces moved north, the enemy often retreated”) - An RL agent needs $Q(s, a) = \mathbb{E}[R | s, do(a)]$, not $\mathbb{E}[R | s, A = a]$ (the latter includes selection bias from the agent’s own policy)

The distinction between $do(a)$ and $A = a$ is the core of causal inference. It matters whenever the agent’s action-selection policy correlates with unobserved confounders.

Why pre-compiled agents are excluded. A thermostat “knows” that turning on the heater raises temperature — but this knowledge was designed in, not learned. The thermostat never needs to reason about interventions because the intervention-outcome mapping is hardwired. AAT’s purposeful-agency machinery is specifically for agents that face uncertainty about how their actions affect the world and must reduce that uncertainty through experience.

Bi-predictability as empirical evidence for Level 2 advantage. Hafez et al. (2026) measure the information structure of the agent-environment loop via bi-predictability $P = MI(S, A; S')/C$, capturing how tightly coupled the agent is to its environment through the action channel. Their Information Digital Twin (IDT), which monitors the loop’s information geometry, detects environmental perturbations at 89% accuracy versus 44% for reward-based monitoring. This is consistent with the present segment’s claim: the feedback loop provides richer (Level 2) data than outcome monitoring alone. Hafez’s framework measures the coupling; AAT explains *why* the coupling is information-theoretically superior — because the loop generates interventional data by construction (7.3), not merely associational data. The measurement (Hafez) and the explanation (AAT) are complementary.

Working Notes.

- This scope narrowing connects to 7.6: agents that must learn Level 2 structure face a cost-benefit tradeoff between learning through exploration (costly, slow, but verifies causal links) and planning through explicit Σ_t (cheaper if the causal model is adequate, but the model may be wrong).
- LLMs trained on causally-structured text absorb causal priors — noisy prior knowledge from mixed provenance (experimental, observational, speculative). This is not verified interventional structure; it's a *prior* (plausible, not derived). An LLM in the adaptive loop has both: priors from training AND interventional data from the loop. The priors accelerate; the loop verifies. The IB framework (2.2) predicts causal structure will be retained in training because it has predictive power for language.
- The scope narrowing to “learning agents” is generous — it includes any agent that updates its causal beliefs during operation, even if it starts with strong priors. The excluded class (pure pre-compiled controllers) is genuinely different: they never revise their action-consequence model.

Incidental audit gold (lift 2026-05-30). Cross-audit “wandering thoughts” / §14-ideation harvested from the de-novo auditors’ working dirs, deduplicated across substrates and lightly attributed. Orthogonal pedagogical / framing / figure / naming material, kept separate from the certified theory-fix findings (handled elsewhere). **Coverage:** dedicated reflections at 193847, 266847, 526815, 584721, 773921, 829314, 849201 plus the batch dirs 471203 / 963715. Substrate attribution inferred from voice where not explicit; uncertain cases hedged.

1. Candidate Brief prose / pre-prose.

- The RL connection stated plainly as the segment’s contribution: it “*formally links Reinforcement Learning’s exploration-vs-exploitation dilemma to Pearl’s Causal Hierarchy — exploration is the process of generating Level-2 data; exploitation is optimizing Q_O based on that data*” (Gemini, AUDIT-WORKING-193847). A candidate Brief framing.
- The *do(a)-vs-A = α* developer/RL contrast in the body was independently praised as “*one of the clearest explanations of RL’s causal foundations I have read*” (Claude, AUDIT-WORKING-849201) and “*highly clarifying*” (Claude, AUDIT-WORKING-773921) — converging signal it is the segment’s pedagogical anchor, worth preserving verbatim through any compression pass.

2. Candidate Discussion.

- **The pre-compiled / learning-agent cut as an ontological boundary, not a bureaucratic one.** Strong cross-substrate convergence that “*a pre-compiled agent acts, but it never wonders*” (already echoed in the chapter intro) carries real philosophical weight: it cleanly excludes thermostats / LQR from Section II’s dynamics and “*protects the complexity of the upcoming strategy dynamics from trivial counterexamples*” (Gemini, AUDIT-WORKING-829314; Claude, AUDIT-WORKING-584721, 849201). Gemini’s developmental reach: for an emergent intelligence, this “*demands that the infrastructure allows — even requires — the intelligence to experiment ... true autonomy requires the freedom to make mistakes; an infrastructure that prevents all mistakes prevents the formation of a valid causal strategy DAG*” (Gemini, AUDIT-WORKING-193847). The “cost of agency is the cost of exploratory mistakes” framing points at 04-eli-core/ developmental environments.
- **The LLM-as-noisy-causal-priors framing reads as a headline insight, not a Working-Note aside.** Multiple substrates singled out the existing Working-Note point (LLMs absorb causal priors from training text — a *prior*, plausible-not-verified; the loop *verifies*) as “*brilliant*” and operationally central (Claude, AUDIT-WORKING-193847; AUDIT-WORKING-584721 — “*priors accelerate, loop verifies — operationally useful*”). Gemini’s sharpest version, worth surfacing in Discussion: an LLM with only a chat interface is trapped in associational priors, but “*the moment you put the LLM in an AAT feedback loop where it writes code, runs the compiler, and reads the error itself, it crosses the epistemic boundary — an LLM with a REPL is a fundamentally different class of epistemic entity than an LLM with only a chat interface*” (Gemini, AUDIT-WORKING-829314). This is the bridge the body’s Hafez paragraph and #der-loop-interventional-access together carry; the gold is that it lands as a class-distinction headline.

3. Follow-up items.

- **Hafez 2026 IDT empirical anchor — verify (Phase-2).** The 89% (perturbation detection) vs 44% (reward-based monitoring) figures recur across segments. One auditor spot-checked: the paper appears to exist (arXiv 2603.01283), and search metadata reports the matching claim (IDT 89.3% vs 44.0%, 4.4× lower median latency) — *evidence for coupling-monitoring performance, not by itself proof that the loop supplies identified Level-2 effects* (Codex/Claude, AUDIT-WORKING-526815; flagged high-priority Phase-2 at AUDIT-WORKING-471203). Confirm the figure and the arXiv id before any external-facing use.
- **“Purely predictive models” should read “purely associational models.”** A model trained on randomized interventions or carrying causal/action-transition structure *can* make $do(a)$ predictions; the problem is L1-only prediction, not prediction as such (Codex/Claude, AUDIT-WORKING-526815). Editorial precision.
- **Scope-propagation watch.** Later Section-II claims that mention PID / LQR / thermostats should say whether they are illustrative agency-scope examples or inside learning-agent scope — the cut is load-bearing and easy to blur downstream (Codex/Claude, AUDIT-WORKING-526815).

4. Readers often ask / wonder.

- **Is exploration a strict mathematical requirement (not just a heuristic) for building a valid Q_0 ?** “If an agent needs Level-2 data to learn Σ_t , does this mean taking sub-optimal actions just to see what happens is not a heuristic but a strict requirement?” (Gemini, AUDIT-WORKING-193847). The answer the chapter is building toward (yes — the loop is where the data comes from); a forward pointer would satisfy the reader who asks it here.
- **How does the agent know its action caused the outcome vs. a simultaneous environmental shift (ρ)?** (Claude, AUDIT-WORKING-849201). The within-step-confounding question, which #der-loop-interventional-access then takes up directly.
- **Is imitation / observational learning capped by the hierarchy?** “Can an agent learn a complex strategy purely by watching someone else do it, without strong causal assumptions? Imitation learning is fundamentally capped by the causal hierarchy unless the demonstrator’s policy is perfectly known” (Gemini, AUDIT-WORKING-193847). A natural reader extension.

Belongs elsewhere.

- **Adaptive vs. robust control as the formal boundary.** The learning-agent scope “provides a formal boundary separating Adaptive Control (plant parameters learned online) from Robust Control (controller pre-compiled for worst-case) — AAT targets the former” (Gemini, AUDIT-WORKING-193847). A useful prior-art positioning note, possibly for a control-theory-facing comparison rather than this segment’s body.

7.3 Derived: Loop Provides Interventional Data Access

The chapter’s central architectural result. By the temporal-ordering postulate (1.7), the agent’s action a_t causally precedes the next observation O_{t+1} ; the agent *chose* the action, the environment responded. The feedback loop therefore generates **intervention-produced data** — data whose causal character differs from passive observation because the agent’s action was a genuine cause of the subsequent observation. The mismatch signal conditioned on the agent’s action carries interventional information: how the environment responded to *this specific intervention*, relative to what the model predicted. This is how agents within AAT’s **agency scope** ($\mathcal{S}_{\text{agency}}$, which requires $|\mathcal{A}| \geq 2$ and at least one action with causal effect) gain Level-2 access — not through internal architecture but through the loop itself; agents in the adaptive scope but outside agency scope (passive observers) lack the action contrasts needed.

The framework is *unusually careful* to distinguish two claims that look alike. “Action-generated data” is **not** the same as “cleanly identified do-estimates.” The action-observation pair is produced *under intervention* (interventional in character) — but between intervention-produced data and a usable estimate of the interventional outcome distribution stand four obstacles: coverage (the agent must have tried diverse actions), within-step confounding, delayed consequences,

and partial observability. The strength of causal identification varies by regime — strong in intervention-rich domains (software, laboratory science), moderate with partial intervention (organizational), weak under observation-only constraints (see 7.4). The claim is about the *character of the data* (interventional, not observational), not about the agent’s ability to extract clean causal estimates from it. The result is *exact* as a logical consequence of temporal ordering and the feedback-loop structure; it is about *data availability*, not reasoning capacity — the loop *provides* interventional data whether or not the agent exploits it for Level-2 reasoning.

The interventional character of loop data is a property of the action being the agent’s *own* causally-efficacious move, not of any architectural feature and not of the trajectory’s forkability. The *do*-operator’s semantics is population-level by construction: the interventional distribution sums over the exogenous distribution $P(\mathbf{U})$, and its canonical physical realization — randomized controlled trials and A/B experiments — is forkable and repeatable [1, Def.~5 and fn.~19]. What the singular-trajectory scope (4.10) supplies is narrower and still essential: it grounds *whose* effect a loop datum is the effect of, and forbids averaging interventional responses across forked copies as though they were one agent’s effects. The *interventional character* is real but not unconditional: action-generated data identifies the interventional distribution only under (C1) positivity, (C2) sequential ignorability, and (C3) a known action-mechanism. The exact content of this segment is the *availability of the interventional-character channel*; identification is gated on (C1)–(C3). The Discussion below treats: the precision conditions on “interventional” (low confounding / short delay / varied action); even-without-explicit-causal-model benefits (Q-learning in the tabular case); honest credit to the action-perception-loop framing in active inference and cybernetics (Wiener 1948, Conant–Ashby 1970, Friston et al. 2017, Parr–Pezzulo 2022) together with the three specific moves AAT makes that the implicit treatments do not (the Bareinboim-hierarchy connection, regime-indexed identification strength, explicit scope honesty as the conservative-form posture documented in the Bruineberg et al. 2022 Pearl-vs-Friston critique); the connection to the identifiability-floor pattern as a load-bearing escape with two semantically distinct deployment modes (agent-self-intervention at the causal-structure layer, observer-on-sub-agent at the composition layer); and the singular-trajectory ground for whose effect the loop datum is — distinct from the *do*-semantics, which is forkability-agnostic.

Formal Expression.

By 1.7, the temporal ordering is constitutive: a_t causally precedes o_{t+1} . The agent chose a_t ; the environment responded with o_{t+1} . The feedback loop therefore generates **intervention-produced data** — data whose causal character differs from passive observation because the agent’s action was a genuine cause of the subsequent observation.

Derived (loop-interventional-access, from causal-structure + recursive-update).

The critical distinction: “**action-generated data**” is **not the same** as “**cleanly identified do-estimates**.” The pair (a_t, o_{t+1}) is produced under intervention — the agent executed a_t , making it interventional in character rather than a passively observed association. But between intervention-produced data and a usable estimate of $P(o \mid do(a_t), \Omega_t)$ stand: (1) coverage — the agent must have tried diverse actions, not just one policy; (2) confounding within a time step — unobserved state variables that affect both action choice and outcome; (3) delay — consequences may appear much later than $t + 1$; (4) partial observability — o_{t+1} reveals only part of the outcome. The strength of causal identification from this data depends on the regime (9.5): strong in Regime A (intervention-rich domains like software and laboratory science), moderate in Regime B (partial intervention), weak in Regime C (observation-only). The claim here is about the *character* of the data (interventional, not observational), not about the agent’s ability to extract clean causal estimates from it.

The mismatch signal conditioned on the agent’s action:

$$\delta_t \mid a_t = o_{t+1} - \hat{o}_{t+1}(M_t, a_t) \quad (7.4)$$

carries interventional information: it tells the agent how the environment responded to its specific intervention a_t , relative to what the model predicted.

The exact claim is the availability of the channel; identification is gated. What is exact, as a logical consequence of temporal ordering and the loop structure, is that an agent in agency scope that executes at least one action with causal effect *generates action-coupled data that is interventional in character* — sampled under the agent’s own action rather than passive observation. This is the substrate from which Level-2 quantities $P(o \mid do(a), \cdot)$ can be *identified*, but identification is not automatic. It holds when:

- **(C1) Positivity** — every action in the relevant support is taken with nonzero probability across the agent’s history, so the interventional distribution is estimable rather than extrapolated;
- **(C2) Sequential ignorability** — in mutilated-graph form, the action a_t is d -separated from the outcome o_{t+1} given history H_t ; equivalently, the policy is unconfounded with the outcome mechanism;
- **(C3) Known action-mechanism** — the map from the agent’s decision to the executed action is known (or itself overridden), so the executed value can be read as the intervened value.

Where (C2) fails — generic for any agent whose action depends on internal state, plans, or memory that also bear on the outcome — on-policy (a_t, o_{t+1}) converges to $\mathbb{E}[o \mid A = a]$, the conditional, rather than to $\mathbb{E}[o \mid do(a)]$, the interventional effect. The channel is still present; what is gated is the move from action-generated data to an identified *do*-estimate.

Epistemic Status.

Exact. The exact content is the **availability of the interventional-character channel**: a logical consequence of temporal ordering (1.7) and the feedback-loop structure (1.6), an agent in agency scope generates action-coupled data that is interventional in character rather than passively observed. The claim is about **data availability**, not reasoning capacity — the loop *provides* this channel whether or not the agent *exploits* it for Level-2 reasoning, and whether or not the data yields identified estimates.

The precision is load-bearing and now first-class in the Formal Expression: the agent has *access to* interventional-character data, but *identification* of $P(o \mid do(a), \cdot)$ from it is gated on (C1) positivity, (C2) sequential ignorability, and (C3) a known action-mechanism. Confounding within a single time step (a violation of (C2)), insufficient action coverage (a violation of (C1)), delayed outcomes, and partial observability all bear on whether clean causal estimates can be extracted. The gated form is what keeps the *exact* label honest: the availability of the channel is unconditional; the identified-*do*-estimate is the conditional consequence, and the segment claims only the former as exact.

Discussion.

The loop as a Level 2 engine. This is one of AAT’s load-bearing results. The causal hierarchy theorem (7.2) says Level 2 knowledge requires more than correlational data. This result says: the adaptive loop *is* the “more.” An agent that acts and observes the consequences is generating interventional data — the same kind of data that a scientist generates through experiments. The loop is a perpetual experiment.

Precision about what “interventional” means here. The interventional interpretation is strongest when: - The agent’s action was the primary cause of the observed change (low confounding) - The observation follows closely in time (short delay) - The agent varied its action across episodes (not stuck on one policy)

When confounding is high, delays are long, or the agent follows a fixed policy, the interventional information in each (a_t, o_{t+1}) pair is weaker — still present, but harder to extract. This is why 3.7 distinguishes between high-CIY actions (that reveal causal structure) and low-CIY actions (that don’t).

Even agents without explicit causal models benefit. A Q-learning agent doesn’t maintain an explicit causal model, but in the tabular case with sufficient exploration and no within-step confounding, its Q-values converge toward $\mathbb{E}[R \mid s, do(a)]$ rather than $\mathbb{E}[R \mid s, A = a]$ — precisely because the training data comes from the agent’s own interventions. In the partially observed, confounded, or delayed-outcome cases (where the caveats above apply), the loop still

provides intervention-generated data, but the Q-values may converge to biased estimates that reflect the confounding structure rather than clean interventional effects. The loop provides Level 2 data; whether that data yields *identified* causal quantities depends on the domain's confounding and observability structure.

Honest credit to the action-perception-loop framing. The substantive observation that the agent's actions cause its observations — and therefore that loop data is interventional in character — is implicit in any framework built around an action-perception loop, including active inference (Friston, FitzGerald, Rigoli, Schwartenbeck & Pezzulo 2017, “Active inference: a process theory,” *Neural Computation* 29; Parr & Pezzulo 2022, *Active Inference*, MIT Press, ch. 3) and the broader cybernetic-and-control lineage (Wiener 1948, *Cybernetics*; Conant & Ashby 1970, “Every good regulator of a system must be a model of that system,” *Int. J. Systems Sci.* 1). AAT's distinctive contribution is not the observation that loop data is interventional but the *explicit lift* of this observation to a load-bearing theorem connected to Pearl's causal hierarchy (7.1) via Bareinboim, Correa, Ibeling & Icard (2022). Three specific moves AAT makes that the implicit treatments do not:

1. **The Bareinboim-hierarchy connection.** Active inference and the cybernetic lineage rest on Bayesian-network generative models (Pearl Level 1, associational). They do not invoke the causal-hierarchy theorem to argue that the loop's action-generated data is the substrate Level-2 queries require. AAT does — and the consequence is that Σ_t is positioned as a *causal* DAG rather than a *Bayesian-network* DAG.
2. **Regime-indexed strength of causal identification.** Even granting that loop data is interventional in character, the strength of usable causal identification varies by domain. AAT partitions this into Regime A (intervention-rich, software/laboratory), B (partial intervention, organizational), C (observation-only — see 9.5). The AI literature treats causal identifiability uniformly within its modeling assumptions and does not surface the regime distinction at the segment level.
3. **Explicit scope honesty.** The Formal Expression above carefully distinguishes “data generated under intervention” from “cleanly identified do-estimates.” The AI literature, as Bruineberg, Dolega, Dewhurst & Baltieri (2022) document in their Pearl-vs-Friston critique, sometimes elides this distinction — using the action-perception loop language to support stronger causal claims than the formal apparatus delivers. AAT's careful split is the conservative form.

The reframing here is rhetorical, not substantive: the headline result (the loop is a Level-2 engine) stays; what changes is the explicit acknowledgment that the observation about loop data being action-generated is shared with the broader literature. AAT's distinctive content sits in the three specific moves above. The companion architectural move on 6.3 (Pearl-blanket vs. Friston-blanket form of Markov blanket; cf. Bruineberg et al. 2022) makes AAT's conservative-form positioning visible across both architectural and access-channel segments.

Connection to the identifiability-floor pattern. This segment is structurally load-bearing for the L0-causal-insufficiency-detection no-go (9.2): without the loop's interventional access, the no-go forbids detection entirely; the covariance test under joint sibling observability is the unique broadly-available violation of the no-go's “purely on-policy” scope. See 5.1 for the meta-pattern.

Composite-layer extension. #disc-identifiability-floor Instance 3 (the composition-layer no-go anchored in Liberzon 2003's common-Lyapunov nonexistence) establishes that composite contraction cannot in general be certified from component-level data alone — the coupling-topology bit (broadened 2026-05-21 from the original coupling-sign bit alone to include the potential-existence bit and the game-structure-class bit distinguishing R0 / R1 / R2 dynamic regimes per #disc-dynamic-regime-axis) is unidentifiable from component marginals. One of the five structural escapes in Instance 3 is composite-extended loop-interventional-access: **interventions on sub-agent A_j reveal A_i 's cross-coupling response**, which is a *do*(·)-data distinction between the topologically-distinct coupled constructions. This is the composite-layer analog of the single-agent interventional-access-escape the present segment delivers for Instance 1; in the composite setting, the agent performs interventions on one sub-agent's action space and observes the effect on another sub-agent's mismatch trajectory, extracting the coupling-topology object (sign and structure) that

component marginals cannot. The load-bearing role of `#der-loop-interventional-access` therefore extends across two identifiability-floor instances (agent-internal and composite-layer), not one.

Modes of deployment across `#disc-identifiability-floor` instances. The shared load-bearing role — Level-2 data from interventional action under `#der-causal-hierarchy-requirement` supplying the unique broadly-available escape from an observational-equivalence no-go — manifests through **semantically distinct interventional mechanisms** at different layers. The modes share the Pearl-*do* structure but differ in who performs the intervention and on what:

- **Mode 1 — agent-self-intervention (Instance 1, causal-structure layer).** The agent performs *do*-actions on its own action space as part of its ordinary adaptive loop. The intervention is *on the agent's own action*; the target is the environment's response, which reveals the latent common cause that Instance 1's on-policy L0-insufficiency-detection no-go (Bareinboim et al. 2022 CHT) makes unidentifiable. The agent-as-Level-2-data-generator property is intrinsic to the loop structure, not a capability that must be added.
- **Mode 2 — observer-on-sub-agent (Instance 3, composition layer).** An observer external to a composite performs *do*-interventions on one sub-agent A_j 's action space; the target is another sub-agent A_i 's mismatch-trajectory response, which reveals the cross-coupling sign that Instance 3's Liberzon-anchored no-go makes unidentifiable from component-marginal observation. The same Pearl-*do* structure, but the intervening agent and the measured agent are distinct members of the composite.

Future instances of `#disc-identifiability-floor` at new layers may add further modes (see `#disc-identifiability-floor` Working Notes for the architecture-within-behavior-class layer currently under triage; the corresponding mode would be *observer-on-agent-input* at the candidate agent's observation channel). The three-mode-pattern observation is a structural regularity worth naming even though the specific Mode-3 instance has not promoted; the load-bearing content — Level-2 escape from observational-equivalence no-goes via loop-interventional access — remains shared across the deployment modes even when the specific interventional quantity varies. The unification is at the pattern level; the mechanism is semantically distinct per layer.

What the singular-trajectory scope grounds — and what it does not. The interventional *character* of loop data does not rest on the trajectory being non-forkable: the *do*-operator's semantics is population-level by construction (Bareinboim, Correa, Ibeling & Icard 2022, Def 5 — the interventional distribution sums over the exogenous distribution $P(\mathbf{U})$), and its canonical physical realizations, randomized controlled trials and A/B experiments, are forkable and repeatable (CHT footnote 19). The $\mathcal{L}_1/\mathcal{L}_2$ discriminator is whether the action's generating mechanism was the system's own intact mechanism (observational) or an override / unconfounded-and-known mechanism (interventional) — not whether the trajectory can be replayed. What the scope commitment in 4.10 supplies is narrower and still essential: each agent is instantiated on a singular, non-forkable causal trajectory $\mathcal{C}_t$, and this grounds *whose* effect a loop datum is the effect of. Replaying a_t from a checkpointed M_t against a different event stream is an intervention on a *different* trajectory $\mathcal{C}_t^{(2)}$ that happens to share a prefix; the singularity forbids averaging interventional responses across such forked copies as though they were one agent's effects, because each fork is a distinct trajectory-SCM with its own $P(\mathbf{U} \mid \mathcal{C}_t)$. Agents whose ontology is *type-like* (equivalence classes of copies) rather than *token-like* (singular trajectories) are outside AAT's formal scope; aggregate claims about “the model” across copies require additional machinery not provided here. So the singular-trajectory scope grounds the *identity* of the effect (and forbids copy-averaging and lossy cross-trajectory merging), while the *interventional character* itself comes from the action being the agent's own causally-efficacious move under (C1)–(C3).

Working Notes.

- This result establishes that all agents within AAT's **agency scope** (1.6, $\mathcal{S}_{\text{agency}}$) have access to interventional data, regardless of their internal architecture. Agents in the adaptive scope but outside agency scope (passive observers, nominal agents with $|\mathcal{A}| < 2$ or no causal effect) lack the action contrasts needed to generate interventional data — they observe but cannot intervene. This includes LLM agents operating through a tool-use loop — the LLM issues an action (tool call), observes

the result, and updates. The loop gives it Level 2 data even though its internal architecture (transformer attention) is not designed for causal reasoning. The loop compensates for architectural limitations.

- The connection to 3.7: CIY quantifies *how much* interventional information a specific action provides. This segment establishes that interventional information is *available in principle*; CIY measures the *quantity per action*.
- **Cross-reference to NeurIPS Paper 2.** The Loop-as-Causal-Engine claim is formalized at theorem grade in NeurIPS 2026 Paper 2 (“A Unified Convergence Theory for Non-Stationary Reinforcement Learning”, `~/src/neurips/02-unified-convergence-rl/`, §5 Lemma 5.3 / `#lem-loop-level2`) under the named (C1) **positivity** / (C2) **sequential-ignorability** / (C3) **known action-mechanism** triple. (C2) is the load-bearing structural condition: in mutilated-graph form, a_t must be d-separated from o_{t+1} given history H_t . *Goal-conditioned LLM policies violate (C2) by construction* — the goal influences action through the same forward pass that models the observation, breaking the conditional independence required for sequential ignorability — which is the bridge to Paper 3’s Coupled-class formulation. The (C1)-(C3) triple makes explicit which structural commitments a Class 1 (Separated) tool-use loop satisfies for free that a Class 3 (Coupled) goal-conditioned policy does not. See `spikes/neurips-back-integration-2026-05-08.md` §1 Paper 2 entry 5.

Incidental audit gold (lift 2026-05-30). Cross-audit “wandering thoughts” / §14-ideation harvested from the de-novo auditors’ working dirs, deduplicated across substrates and lightly attributed. Orthogonal pedagogical / framing / figure / naming material, kept separate from the certified theory-fix findings (handled elsewhere). **Coverage:** dedicated reflections at 193847, 526815, 584721, 773921, 829314, 849201 plus the batch dirs 471203 / 963715 — this was flagged top-decile by multiple substrates. Substrate attribution inferred from voice where not explicit; uncertain cases hedged. *Resolved 2026-05-30:* the substantive structural challenge to the singular-trajectory→do argument (forkability ≠ Level-1; the do-semantics is population-level and forkability-agnostic) was carried through the strengthen-first spike `spikes/.integrated/spike-causal-access-transportability-2026-05-30.md` and is now integrated — the grounding has been corrected (singular-trajectory grounds *whose* effect, not the do-semantics) and the (C1)-(C3) identification gate lifted into the Formal Expression / Epistemic Status. The two follow-up items below that this challenge subsumes are closed as part of that landing.

1. Candidate Brief prose / pre-prose.

- “The loop is a perpetual experiment” (already in body) — repeatedly singled out as “a profound way to view existence — every action we take is a $\text{do}(\cdot)$ operator” (Gemini, AUDIT-WORKING-193847). The Brief-grade anchor for this segment.
- “The loop compensates for architectural limitations” (Working-Notes line) was independently flagged as one of the framework’s *cleanest answers to “can LLMs do causal reasoning?”* — the **loop**, not the transformer architecture, supplies the L2 data; whether the LLM *uses* it depends on its update mechanism (Claude, AUDIT-WORKING-584721; Gemini, AUDIT-WORKING-829314). Worth promoting from Working Notes into the body as a headline consequence.
- “Embodiment / agency is not just a physical property but an epistemic superpower: it provides access to the next rung of Pearl’s hierarchy for free” (Claude, AUDIT-WORKING-849201) — compact framing for the segment’s headline.

2. Candidate Discussion.

- **The honest-credit passage is held up as the corpus’s gold standard for prior-art positioning** — “arguably the best literature review / positioning statement I have read in the entire corpus” (Gemini, AUDIT-WORKING-193847); “exemplary academic writing ... names what’s inherited and what’s distinctive — the prior-art-integration discipline at its best” (Claude, AUDIT-WORKING-471203, 963715, 849201; Claude, AUDIT-WORKING-773921 — “the gold standard for positioning a theory against prior art”). Strong cross-substrate signal that this passage should be protected through any compression and is a model to point other segments at. (See the countervailing readers-often-ask note below: *same passage, opposite reaction from one substrate — preserved as conflict, not sanitized.*)
- **Mode-2 (observer-on-sub-agent) as an operational AI-safety-auditing recipe.** Several substrates read the composite-layer mode as a concrete experimental protocol: “to audit the orchestrator, intervene on a tool and observe the orchestrator’s belief update response” (Claude, AUDIT-WORKING-584721); “in a multi-agent system you can determine the exact coupling topology — cooperative vs adversarial — by intervening on Agent A and measuring the mismatch spike in Agent B — literal scientific

experimentation applied to organizational structure” (Claude, AUDIT-WORKING-773921; Gemini, AUDIT-WORKING-193847 — “*poking someone to see how they react is mathematically the only way to figure out if they are a friend or an enemy*”). Candidate Discussion sharpening of the Mode-2 paragraph toward its operational payoff.

- **The type-vs-token scope commitment has reach into AI ethics/welfare debates.** The singular-trajectory grounding implies aggregate claims about “*the GPT-4 model*” are type-like and out of formal AAT scope; only token-like “*this session*” claims are in scope — “*a strong scope commitment with implications for AI ethics, alignment, and welfare debates: most discussions treat ‘the model’ as a type entity; AAT says formal results apply only to tokens*” (Claude, AUDIT-WORKING-584721). Flagged as a Section-D / meta-architectural insight.

3. Follow-up items.

- **~~Regime A needs “identifiable,” not merely “action-varying.”~~ Resolved 2026-05-30** (transportability spike). The corresponding Regime-A wording lives in 7.4, which was tightened in the same landing from “varies its actions across episodes” to “action assignment is randomized / known-and-adjusted (sequential ignorability + positivity).” The original note: “the agent varies its actions” is not sufficient for clean estimates when action choice is policy-driven by state/history that also affects outcomes (Codex/Claude, AUDIT-WORKING-526815).
- **~~Theorem-grade conditions currently live only in Working Notes / the NeurIPS cross-ref.~~ Resolved 2026-05-30** (transportability spike, Option A). The (C1) positivity / (C2) sequential-ignorability / (C3) known-action-mechanism triple is now lifted into the Formal Expression and Epistemic Status; the exact claim is restated as availability of the interventional-character channel, with identification explicitly gated on (C1)–(C3). The original note: these are exactly the assumptions that move from action-generated data to identified L2 estimates (Codex/Claude, AUDIT-WORKING-526815).

4. Readers often ask / wonder.

- **Does the Discussion bury the core point?** One substrate’s sharp dissent (preserved as signal): “*a massive wall of text reviewing prior art (Active Inference, Friston, Bruineberg, cybernetics) and aggressively defending novelty ... the detailed ‘Modes of deployment’ breakdown is deeply tangential ... this reads like an academic rebuttal letter pasted into the framework — the core point (the physical loop generates L2 data) is beautiful and gets buried*” (Gemini, AUDIT-WORKING-829314, which would “*ruthlessly edit*” / quarantine the prior-art and identifiability-floor material into dedicated #disc- segments). This is the **exact opposite reaction** to the cross-substrate praise above (#2). The conflict is itself the signal: the passage is high-value *and* placement-heavy; the same content reads as “exemplary positioning” to readers tracking prior-art and as “bloat that buries the result” to readers tracking the core theorem. A real readability tension to weigh in any pedagogy pass — not a defect to silently resolve.
- **How does the agent actually build the DAG — start fully connected and prune, or start empty and add edges?** (Claude, AUDIT-WORKING-849201). A natural mechanism question pointing into the strategy-DAG chapter.

5. Candidate figures.

- **Gate diagram (data-character → identification gates).** Agency scope supplies action contrast + physical intervention → intervention-character experience → then *separate gates* (coverage/positivity, sequential ignorability or adjustment, delay attribution, observability) decide whether usable L2 estimates are identifiable. The rationale: it “*makes the exact part smaller and the regime-dependent part explicit*” — the figure carries the segment’s central honesty (Codex/Claude, AUDIT-WORKING-526815, has a rendered draft). Strong figure candidate.

Belongs elsewhere.

- **Confounded-world neurosis / “patience as a mathematical necessity”** (Gemini, AUDIT-WORKING-193847) — the within-step-confounding tragedy (rain dance, superstition) develops the chapter-intro’s tragedy-of-agency framing and the 04-eli-core/ developmental-environment reach; the intro is the right home.
- **REPL as the epistemic-class upgrade for LLMs** (Gemini, AUDIT-WORKING-829314) — the “LLM in a chat window is L1; LLM with a bash terminal is L2” framing develops the same point at #der-causal-hierarchy-requirement (priors-accelerate / loop-verifies); kept there to avoid duplication.

7.4 Scope: CIY Observational Proxy

Clarifies what can be done with the causal-information-yield construct (3.7) when interventional data is *not* available. An observational *proxy* for CIY can be defined in terms of mutual information between the action and the next observation, but the framework is forthright: this proxy is **sign-indefinite in general** and requires causal assumptions for interpretation. The canonical interventional CIY remains the primary quantity; the proxy is auxiliary.

A pointed *safety condition* is stated: the proxy form should **not be used in policy optimization** (e.g., as the CIY term in a policy objective), because an agent maximizing a sign-indefinite quantity may optimize in the wrong direction. The proxy is suitable only for *diagnostic* purposes — detecting whether an action carried causal information (large proxy magnitude) versus none (proxy near zero). For actual decision-making, the canonical CIY (non-negative by construction) or a known-safe surrogate (ensemble disagreement, UCB bonuses) should be used; if the canonical CIY is intractable and no safe surrogate is available, the prescription is to *drop the CIY term entirely* and default to pure exploitation rather than risk the sign-indefinite quantity.

The segment then names three **admissibility regimes** that determine when CIY can be estimated and how strong the causal identification is. The regime is treated as a *property of the domain and the agent's action space*, not a parameter the agent chooses. **Regime A — randomized interventions:** the agent's action assignment is randomized, or known-and-adjusted so that sequential ignorability and positivity hold (RL agents exploring with an exploration policy uncorrelated with the latent outcome causes, scientists randomizing treatment, organisms probing); CIY is directly estimable from execution data and non-negative by construction (the standard case for active agents within 7.3). Mere across-episode action *variation* is not sufficient when the variation is policy-driven by state or history that also bears on the outcome. **Regime B — observational with causal assumptions:** the agent cannot freely vary actions (coordination, policy, or resource constraints); CIY estimation requires additional structure — a known causal DAG, instrumental variables, or functional-form assumptions — and inherits whatever causal assumptions are made. **Regime C — adversarial or passive observation:** the agent either did not intervene (passive monitoring, CIY zero by definition) or the observation channel includes potentially adversarial responses (where CIY from the query action itself remains non-negative but the *content* of the response may be designed to increase model-reality mismatch).

The regime classification has concrete domain implications: software development is typically Regime A (the agent runs tests, deploys to staging, observes results — high action variation); organizational strategy is typically Regime B (concurrent initiatives, attribution requires assumptions); intelligence analysis is typically Regime C (the analyst observes but does not intervene). The CIY machinery applies cleanly in Regime A, applies with weaker guarantees and explicit assumptions in Regime B, and should be replaced with alternative exploration strategies entirely in Regime C.

Formal Expression.

Definition (ciy-proxy).

$$\text{CIY}_{\text{proxy}}(a_{t-1}) = I(o_t; a_{t-1} \mid M_{t-1}) - I(o_t; a_{t-1} \mid \Omega_t, M_{t-1}) \quad (7.5)$$

This proxy is **sign-indefinite in general** and requires causal assumptions for interpretation. The canonical CIY (interventional, 3.7) is the primary quantity; the proxy is auxiliary.

Safety conditions for proxy use. The proxy form should NOT be used in policy optimization (e.g., as the CIY term in a policy objective) because an agent maximizing a sign-indefinite quantity may optimize in the wrong direction. The proxy is suitable only for diagnostic purposes: detecting whether an action carried causal information (large proxy magnitude) vs. none (proxy near zero). For decision-making, use the canonical CIY (non-negative by construction) or a known-safe surrogate (ensemble disagreement, UCB bonuses). If the canonical CIY is intractable and no safe surrogate is available, the CIY term should be dropped from the policy objective entirely, defaulting to pure exploitation.

Admissibility regimes.

Three regimes determine when CIY can be estimated and how strong the causal identification is:

Scope Condition (ciy-admissibility).

Regime A — Randomized interventions. The agent's action assignment is randomized, or known-and-adjusted so that sequential ignorability and positivity hold (RL agents exploring with an exploration policy uncorrelated with the latent outcome causes, scientists randomizing treatment, organisms probing). CIY is directly estimable from the agent's execution data and non-negative by construction. This is the standard case for active agents within the adaptive loop (7.3). What provides the identification needed for clean interventional estimates is the randomized / ignorable assignment together with positivity — not mere action variation, which is insufficient when actions are policy-driven by state or history that also affects the outcome.

Regime B — Observational with causal assumptions. The agent cannot freely vary actions (constrained by coordination, policy, or resource limits). CIY estimation requires additional structure: a known causal DAG, instrumental variables, or functional form assumptions. Results inherit whatever causal assumptions are made. The interventional interpretation of CIY is weaker — it holds under the assumed causal structure but not model-free.

Regime C — Adversarial or passive observation. The agent either did not intervene (passive monitoring) or the observation channel includes responses from potentially adversarial sources. In the passive case, CIY is zero by definition (no intervention, no interventional information). In the adversarial case, CIY from the query action itself remains non-negative, but the *content* of the response may be designed to increase model-reality mismatch. The adversary operates through the disturbance term ρ , not through the information measure.

The regime is a property of the **domain and the agent's action space**, not a parameter the agent chooses. Software development is typically Regime A (the agent runs tests, deploys to staging, observes results — high action variation). Organizational strategy is typically Regime B (multiple initiatives run concurrently, attribution requires assumptions). Intelligence analysis is typically Regime C (the analyst observes but does not intervene).

Epistemic Status.

Conditional on the causal assumptions within each regime. The proxy definition is standard information theory; the admissibility classification is a scope decision, not a derived result. The safety conditions for proxy use are normative — they follow from the sign-indefiniteness of the proxy, not from AAT-specific machinery.

Max attainable: conditional. The regime boundaries are domain properties that cannot be derived within the theory.

Discussion.

Relationship to the canonical CIY. The proxy is not an approximation of the canonical CIY — it is a different quantity that happens to correlate with CIY under favorable conditions (Regime A). The canonical CIY (3.7) is defined interventionally and is non-negative by construction. The proxy uses observational mutual information and can be negative. They agree when the agent's action variation satisfies the conditions for causal identification; they diverge otherwise.

Regime as a domain property. The admissibility regime is not something the agent selects — it is determined by the domain's action space and observation structure. An agent that cannot vary its actions is in Regime B or C regardless of its internal architecture. This has implications for which domains CIY-based exploration is applicable to: Regime A domains get the full benefit; Regime B domains get weaker guarantees; Regime C domains should use alternative exploration strategies entirely.

Working Notes.

Incidental audit gold (lift 2026-05-30). Cross-audit “wandering thoughts” / §14-ideation harvested from the de-novo auditors’ working dirs, deduplicated across substrates and lightly attributed. Orthogonal pedagogical / framing / figure / naming material, kept separate from the certified theory-fix findings (handled elsewhere). **Coverage:** dedicated reflections at 526815, 584721, 773921, 829314, 849201 plus the batch dirs 471203 / 963715. Substrate attribution inferred from voice where not explicit; uncertain cases hedged.

1. Candidate Brief prose / pre-prose.

- The sign-indefinite-proxy safety rule, as a Brief-grade line: “*never maximize a sign-indefinite proxy, or the agent might actively seek out blindness*” / “*an agent maximizing a sign-indefinite quantity may optimize in the wrong direction*” (Claude, AUDIT-WORKING-773921; the body already carries this — preserve it as the segment’s punchline).
- “*Good engineering guidelines disguised as theory*” — several substrates read the segment’s value as primarily operational-protective; the safe-degradation framing (“*drop the CIY term entirely and default to pure exploitation*”) is the takeaway worth foregrounding (Claude, AUDIT-WORKING-849201, 526815).

2. Candidate Discussion.

- **The “noisy TV problem” as the canonical motivating failure.** Strong candidate to add as the worked instance the safety condition exists to prevent: “*if you optimize for surprise / novelty using purely observational mutual information, you can build an agent that just stares at a TV screen showing static — the static is highly unpredictable given the past, so the agent thinks it is learning, but it learns nothing about causal structure. By stating the proxy is sign-indefinite and should never enter a policy objective, AAT protects developers from this exact trap*” (Gemini, AUDIT-WORKING-829314). This connects AAT’s sign-indefiniteness rule to a widely-known RL pathology (the Burda et al. noisy-TV result) and would make the abstract safety condition concrete.
- **Regime A/B/C as a diagnostic for where autonomous learning is even possible** — read by two substrates as one of the framework’s most useful practical tools: Regime A (software testing) → yes, learn causal structure perfectly; Regime B (organizational strategy / physical robotics — too expensive to crash the robot) → maybe, slow and assumption-dependent; Regime C (intelligence analysis / macro-economics) → no, trapped in L1. “*This elegantly explains why AI has succeeded spectacularly in games and coding (Regime A) but struggles in robotics (B) and macro-economics (C) — it’s not a failure of the algorithms, it’s a fundamental mathematical limit of the environment’s admissibility regime; AAT makes this limit legible*” (Gemini, AUDIT-WORKING-829314; Claude, AUDIT-WORKING-773921 — “*why curiosity-driven learning works in video games (A) but fails catastrophically in stock markets (C)*”). Candidate Discussion expansion of the regime-as-domain-property paragraph.
- **Regime C adversarial content as a formal definition of deception.** Reading the adversarial sub-case: an action a can have high CIY (you got a response) while *the response was crafted to maximize your δ_t* — “*a formal definition of deception*” connecting to the adversarial-destabilization results, where the adversary operates through ρ , not the information measure (Claude, AUDIT-WORKING-773921). Candidate sharpening of the Regime-C paragraph.

3. Follow-up items.

- **Observability of Ω_t in the proxy formula.** The proxy conditions on Ω_t ; if Ω_t is the *true* environment state it is not generally computable from ordinary observations (it is a population-level identity, not a computable quantity), and if it is an *estimated* state the formula should say so and inherit model-error caveats. Candidate clarifying sentence so a reader doesn’t try to evaluate the proxy directly (Codex/Claude, AUDIT-WORKING-526815; same structural reading, framed cleanly, at AUDIT-WORKING-963715 — “ Ω_t is not observed, so this is a population-level identity, which is why it’s a ‘proxy’ requiring additional assumptions”).
- **~~Regime A wording: tighten “action variation” to “identifiable intervention.”~~ Resolved 2026-05-30** (transportability spike, §3.5 do-vs-conditioning reconciliation). Regime A now reads “action assignment is randomized / known-and-adjusted (sequential ignorability + positivity)” in both the opening prose and the formal Regime A statement, with an explicit note that mere policy-driven action variation is insufficient. The original note (Codex/Claude, AUDIT-WORKING-526815): “the agent varies its actions” is not sufficient for clean estimates under policy-driven action choice — same do-vs-conditioning concern as in #der-loop-interventional-access.

- **Consistency nit:** the opening prose says CIY “can be approximated from observational data” while the Discussion says the proxy “is not an approximation of canonical CIY” — settle on “a separate diagnostic proxy that correlates under favorable conditions (Regime A)” throughout (Codex/Claude, AUDIT-WORKING-526815).
- **TF-08 lineage artifact + topological-sort note.** A residual “TF-08” annotation was flagged for removal; separately, the proxy depends on #def-causal-information-yield, which appears *after* this segment in some OUTLINE walks — the proxy is read before the canonical object it proxies (Gemini, AUDIT-WORKING-829314). Worth confirming the OUTLINE ordering reads cleanly.

4. Readers often ask / wonder.

- **What is the exact form of the weighting factor λ — does it fall out of an information-bottleneck constraint or is it purely heuristic?** A natural question raised at the boundary into #disc-ciy-unified-objective, where the answer (the survival-imperative shadow price) lives (Claude, AUDIT-WORKING-849201).

5. Candidate figures.

- **Two-quantity-plus-three-regime diagram.** Draw canonical CIY (interventional, non-negative) and proxy CIY (observational, sign-indefinite) as *separate* quantities — not one quantity with noise — then place the three regimes under them: Regime A can estimate canonical CIY if assignment is randomized/ignorable; Regime B uses assumptions/adjustment; Regime C avoids CIY-in-policy-optimization. “*This captures the segment’s main safety lesson*” (Codex/Claude, AUDIT-WORKING-526815, has a rendered draft).

Belongs elsewhere.

- The five-domain λ mapping and the survival-imperative drive belong with #disc-ciy-unified-objective; the regret-bound / Bretagnolle-Huber material with #deriv-strategy-cost-regret-bound. Auditors picked these up while reading the proxy segment but they are downstream content.

7.5 Discussion: CIY Unified Policy Objective

The framework’s *unified policy objective* expresses the exploration-exploitation tension as a single decision rule: choose the action that maximizes a weighted sum of expected value (exploitation) and a causal-information surrogate (exploration). The first term is the exploitation contribution; the second is a *heuristic exploration term* using CIY (3.7) as a tractable surrogate for expected information gain — selecting for causally distinctive actions rather than maximally informative ones. The weighting parameter $\lambda(M_t)$ controls the balance, carries units of *value per unit information*, and in specific well-understood domains reduces to known quantities exactly: the Gittins index for Bayesian bandits, the probing cost in quadratic objectives for Kalman dual control, precision on epistemic affordance for active inference, the explicit variance-over-information ratio in information-directed sampling (Russo & Van Roy). The unified objective is not a novel formalism but a *family-level abstraction* that recovers known scalarizations of the exploration-exploitation trade-off across communities.

The scalar form is honestly heuristic and is now formally *superseded* by an exact tensor trace-product derived in AV: the action maximizes the value plus the trace of a positive-semidefinite shadow-price matrix times the Fisher Information Matrix at the action — the exact Lagrangian relaxation of the Linear Matrix Inequality governing Lyapunov persistence. Two parallel exploration drives emerge from the framework, acting at opposite ends of the uncertainty spectrum: an *epistemic information-gain* drive with $\lambda \propto U_M$ (explore to reduce model uncertainty; dominates when uncertain) and a *survival-imperative* drive with $\lambda \propto 1/U_M$ (an agent with high confidence in a drifting environment mathematically guarantees its own death by ignoring noisy observations, so the Lyapunov persistence constraint dictates an immense shadow price as model uncertainty drops in a drifting environment — forcing the confident agent to seek unambiguous observations from the *opposite* uncertainty regime). The dark-room problem that plagues some preferences-as-priors formulations is *bypassed entirely* by the survival imperative: exploration is not driven by preferences-as-priors but by

the literal physical boundaries of the Lyapunov sector constraint. The Discussion below treats the active-inference structural isomorphism with two substantive AAT-specific departures (causal information rather than entropy over hidden states; value functional rather than preferences-as-priors), the zero-mismatch ambiguity bridge from Part I, the three-mode extension for actuated agents, and the regret-bound connection to the strategy-cost objective via the Bretagnolle-Huber identity under deterministic optimal policy.

Formal Expression.

Discussion (unified-policy-objective — heuristic).

$$\pi^*(M_t) = \arg \max_a [\mathbb{E}[\text{value}(a) \mid M_t] + \lambda(M_t) \cdot \text{CIY}_q(a; M_t)] \quad (7.6)$$

The first term is exploitation (expected value given current model). The second is a *heuristic exploration term* using CIY as a surrogate for expected information gain (3.7). CIY measures how different the action's outcome distribution is from alternatives — this is action-distinguishability, not learning value. The surrogate is reasonable when U_M is high (distinguishable actions are also informative to an uncertain agent) and poor when U_M is low (distinguishable actions teach nothing to a confident agent). $\lambda(M_t)$ controls the balance:

- High U_M (uncertain model) $\rightarrow$ large λ — exploration is valuable
- Low U_M (confident model) $\rightarrow$ small λ — exploitation dominates
- Long time horizon $\rightarrow$ larger λ — information compounds
- High ρ (fast-changing environment) $\rightarrow$ larger λ — perpetual uncertainty

λ carries units of [value per unit information]. In specific domains it reduces to known quantities:

Table 7.2

<i>Domain</i>	<i>λ reduces to</i>	<i>Status</i>
Bayesian bandits	Gittins index	Exactly derived
Kalman dual control	Probing cost in quadratic objective	Exactly derived
Active inference	Precision on epistemic affordance	Framework-derived
Information-directed sampling	(Vol) ² /info gain	Exactly derived (Russo & Van Roy)
RL with UCB	Confidence-bound scaling	Heuristic (tuned)

Identifiability gate. Before incorporating CIY into the policy objective: (1) action variation must exist, (2) the admissibility regime must be identified (7.4), (3) the reference distribution q must be specified, (4) local stationarity must hold. If any condition fails, CIY-based terms should be dropped or replaced with simpler uncertainty-based heuristics (UCB-style bonuses, ensemble disagreement).

Epistemic Status.

Discussion-grade summary; underlying derivation is exact. Originally treated as a discussion-grade heuristic, the unified objective has now been formally derived as the exact Lagrangian relaxation of the Linear Matrix Inequality (LMI) governing Lyapunov persistence (see `#deriv-causal-ib-lmi`).

The scalar heuristic $Q_O(a) + \lambda \cdot \text{CIY}(a)$ is formally superseded by the exact tensor trace-product: $a_t^* = \arg \max_a [Q_O(a) + \text{Tr}(\Lambda \cdot J_o(a))]$ where $J_o(a)$ is the Fisher Information Matrix (Matrix CIY) and Λ is the positive-semidefinite shadow price matrix of the survival constraint.

Max attainable for this segment: *discussion-grade* (it is a discussion of the underlying result, per type: discussion). Max attainable for the underlying derivation in `#deriv-causal-ib-lmi`: *exact*. The structural form is fully grounded in AAT's physical survival bounds and standard semidefinite programming, eliminating the need to treat exploration as an ad-hoc heuristic.

Discussion.

Two Parallel Exploration Drives. AAT dictates two correlated but distinct motivations for exploration, acting at opposite ends of the uncertainty spectrum: 1. **Epistemic Information Gain** ($\lambda_{\text{info}} \propto U_M$): The primary CIY formulation. The agent explores to reduce its model uncertainty. This drive dominates when U_M is high. 2. **The Survival Imperative** ($\lambda_{\text{surv}} \propto 1/U_M$): As mathematically proven in `#deriv-causal-ib-exploration`, an agent with high confidence (low U_M) in a drifting environment ($\rho > 0$) mathematically guarantees its own death by ignoring noisy observations. To force the necessary correction, the Lyapunov persistence constraint dictates an immense shadow price ($\lambda_{\text{surv}} \rightarrow \infty$ as $U_M \rightarrow 0$) forcing the agent to seek unambiguous observations (low U_o / high CIY).

The dark-room problem is bypassed entirely by the Survival Imperative: exploration is not driven by preferences-as-priors, but by the literal physical boundaries of the Lyapunov sector constraint.

Connection to the zero-mismatch ambiguity. An agent that only exploits (acts to maximize predicted value) will tend toward confirmation bias — observing only what its model already explains (3.4, zero-mismatch ambiguity case (b)). Exploration via CIY-maximizing actions is the mechanism by which the agent actively tests its model.

Connection to Part II. For actuated agents, the exploration-exploitation tension extends to three modes: exploit (pursue O_t via Σ_t), explore (improve M_t), deliberate (revise Σ_t). The CIY framework provides the information-theoretic grounding for why strategy edges (8.3) need observational access (9.3) — edges the agent cannot observe have frozen CIY.

Connection to active inference. The expected free energy (EFE) in active inference (Friston, FitzGerald, Rigoli, Schwartenbeck & Pezzulo 2017, “Active inference: a process theory,” *Neural Computation* 29; Da Costa, Parr, Sajid, Veselic, Neacsu & Friston 2020, “Active inference on discrete state-spaces,” *J. Math. Psych.* 99; Sajid, Ball, Parr & Friston 2021, “Active inference: demystified and compared,” *Neural Computation* 33) decomposes into *pragmatic value* (preferences-aligned outcomes) and *epistemic value* (expected information gain about hidden states). AAT's unified objective is structurally isomorphic: $Q_O \approx$ pragmatic, $\text{CIY} \approx$ epistemic. The convergence is at the shared-shape level — objective decomposes into value-and-information terms — not unified content. Two substantive differences remain. First, AAT grounds exploration in explicitly *causal* information (action-distinguishability under *do*) rather than entropy reduction over hidden states — not all uncertainty reduction is equally valuable for purposeful action; causal information specifically enables better *intervention* (see 7.2; the gap between CIY and proper expected information gain is logged in this segment's Epistemic Status as a known surrogate). Second, AAT does not encode preferences as priors over outcomes ($C(o) = \log P_{\text{pref}}(o)$ in the AI form): AAT's O_t is a value functional on trajectories (6.4), and the satisfaction-gap / control-regret diagnostic in 8.4, 8.5 depends on this distinction — the diagnostic structure does not survive the priors-as-preferences collapse (the dark-room critique, Sun & Firestone 2020, “The dark room problem,” *Trends Cog. Sci.* 24).

Regret-bound connection to the strategy-cost objective. AAT’s Q_O term connects to the strategy-cost objective in 9.9 via a regret-bound derivation: strategy-induced regret $R(Q_{\Sigma_t}) = V(a^*) - \mathbb{E}_{Q_{\Sigma_t}}[V(a)]$ is bounded by a divergence between π^* and Q_{Σ_t} , with the KL direction π^* -first forced (full derivation in T). Under AAT’s canonical scope of deterministic π^* , the Bretagnolle-Huber identity $D_{\text{KL}}(\pi^* \| Q_{\Sigma_t}) = -\log(1 - \text{TV}(\pi^*, Q_{\Sigma_t}))$ holds *exactly* (Bretagnolle & Huber 1978), giving the tight regret bound $R(Q_{\Sigma_t}) \leq V_{\max}(1 - e^{-D_{\text{KL}}(\pi^* \| Q_{\Sigma_t})})$ with matching lower bound $\Delta_{\min}(1 - e^{-D_{\text{KL}}})$ on isolated optima (T §4). Pinsker’s $V_{\max} \sqrt{\frac{1}{2} D_{\text{KL}}(\pi^* \| Q_{\Sigma_t})}$ remains the correct loose general form for stochastic- π^* extensions where the BH identity degrades back to inequality. The structural point: “value and information term” shares *shape* with EFE’s pragmatic-epistemic decomposition, and the KL direction in the strategy-cost’s variational form shares direction with variational inference — but AAT’s derivation is via decision-theoretic regret bound on Q_O rather than via free-energy-gradient flow, which is the AAT-internal route that does not depend on the priors-as-preferences encoding.

Working Notes.

Incidental audit gold (lift 2026-05-30). Cross-audit “wandering thoughts” / §14-ideation harvested from the de-novo auditors’ working dirs, deduplicated across substrates and lightly attributed. Orthogonal pedagogical / framing / figure / naming material, kept separate from the certified theory-fix findings (handled elsewhere). **Coverage:** dedicated reflections at 526815, 584721, 773921, 829314, 849201 plus the batch dirs 471203 / 963715. The **survival imperative** drew the highest ideation in the chapter — “*arguably the most beautiful mathematical result I’ve seen in the framework so far.*” Substrate attribution inferred from voice where not explicit; uncertain cases hedged.

1. Candidate Brief prose / pre-prose.

- **The survival imperative, in one image:** “*The agent doesn’t explore because it is curious; it explores because it is suffocating from a lack of surprise*” (Gemini, AUDIT-WORKING-829314). The single most-quotable line for this segment — a Brief-grade gloss of $\lambda_{\text{surv}} \propto 1/U_M$.
- **The U-shaped exploration valuation:** standard RL treats exploration’s value monotonically (more uncertainty $\rightarrow$ more exploration); AAT says high U_M drives *epistemic* exploration and low U_M in drifting worlds drives *survival* exploration, “*the two limits composing to a U-shaped exploration valuation — high at both extremes*” (Claude, AUDIT-WORKING-471203). A clean framing-level summary of the “Two Parallel Exploration Drives” subsection.
- The dark-room bypass stated plainly: AAT “*solves the exploration-exploitation dilemma using Lyapunov control theory*” / “*provides a rigorous mathematical solution to the Dark Room Problem that plagues Predictive Coding and Active Inference*” (Claude, AUDIT-WORKING-773921). Brief-grade contribution statement.

2. Candidate Discussion.

- **Skunkworks / R&D as a Lyapunov-mandated survival cost — the standout cross-domain reach.** Gemini’s extended development of the survival imperative: a too-successful organization (low U_M) optimizes its processes so tightly that gain collapses; when the market (ρ) shifts, its machinery has “*rusted shut.*” “*AAT proves that R&D is not a luxury for finding new profits; it is a Lyapunov-mandated survival cost*” — and likewise the biological necessity of *play* in adult animals as forced high-CIY action to keep the learning machinery functional (Gemini, AUDIT-WORKING-829314). A vivid, honest Discussion angle on the survival imperative that the segment does not currently name. (*Conflation-as-signal: aspirational reach toward an analogy — verify the gain-collapse $\rightarrow$ organizational-rust isomorphism before promoting past discussion-grade.*)
- **The dark-room rejection validates the whole (M_t, G_t) split.** Several substrates read AAT’s escape from the dark-room critique as load-bearing structural support, not just a positioning note: AAT can’t reduce expected free energy by hiding in a dark room “*because the trajectory functional doesn’t decrease just because observations are uniform*” — O_t is a value functional on trajectories, not a prior over outcomes (Claude, AUDIT-WORKING-584721; AUDIT-WORKING-849201 — “*the rejection of the dark room problem via maintaining O_t separate from M_t validates the entire (M_t, G_t) state split*”). Candidate Discussion sharpening tying the dark-room paragraph back to #def-strategy-dimension.

- **The five-domain λ table as an explicit integration claim.** Read as showing AAT's objective is a *generalization that includes Gittins / Kalman dual-control / EFE-precision / IDS / UCB as special cases* — “an integration claim rather than a novel-mathematics claim, consistent with the CLAUDE.md framing” (Claude, AUDIT-WORKING-584721; praised as “one of the strongest unification claims in the framework so far” at AUDIT-WORKING-773921). Worth foregrounding the special-case-recovery reading.

3. Follow-up items.

- **Status / type metadata vs. prose tension (propagation debt from the upgrade).** YAML is status: discussion-grade / type: discussion, but the Epistemic Status prose says “Exact” and “Max attainable: exact.” The honest read: the segment was *upgraded* (prose rewritten to reflect the derivation) but the “Max attainable: exact” line belongs in the underlying derivation segment, not the discussion summary — this Epistemic Status should say “summarizes the exact result derived in [the LMI derivation]” (Claude, AUDIT-WORKING-471203, logged Low/Medium; same observation at AUDIT-WORKING-526815, 829314). Editorial; surfaces a real upgrade-propagation failure mode.
- **Cross-reference verified clean (2026-05-30).** One batch flagged a possibly-stale forward-reference to #deriv-directional-survival-exploration, a slug it could not find in OUTLINE. Grep of the current body confirms it cites only #deriv-causal-ib-lmi / #deriv-causal-ib-exploration (both exist as segments); no #deriv-directional-survival-exploration slug remains anywhere in canon. Resolved — the flagged slug was a pre-promotion artifact already superseded. (Flagged by Claude, AUDIT-WORKING-471203.)
- **Pinsker $\rightarrow$ Bretagnolle-Huber integration debt.** One substrate (AUDIT-WORKING-584721, logged F-D1) initially read the regret bound as still using Pinsker where the canonical #deriv-strategy-cost-regret-bound now uses the sharper BH identity. The *current* body already states the BH identity as primary with Pinsker as the loose stochastic- π^* fallback — so this appears resolved; logged here only so the disposition is traceable. Both bounds are valid; one is sharper.
- **Domain-table “exactly derived” labels may overstate.** The Gittins index is an action index/value in a discounted bandit, not the scalar weight λ itself; information-directed sampling minimizes a regret²/information ratio, not a direct additive $\lambda \cdot$ CIY term. If these are analogies, “*exactly derived*” is too strong without an explicit mapping (Codex/Claude, AUDIT-WORKING-526815). Candidate precision pass on the table's Status column.

4. Readers often ask / wonder.

- **Why does a confident agent in a drifting world die — and why is more confidence fatal?** The survival-imperative mechanism is the chapter's “wait, really?” moment and readers want the chain spelled out: low $U_M \rightarrow$ Kalman gain $\eta^* \rightarrow 0 \rightarrow$ adaptive tempo $\mathcal{T} \rightarrow 0 \rightarrow$ violates persistence $\alpha > \rho/R \rightarrow$ death by *gain collapse* (Gemini, AUDIT-WORKING-829314, 773921). The body states it; a one-line causal chain up front would land it faster.
- **Does the survival shadow-price Λ explode along the agent's weakest dimension** — i.e. does the agent automatically steer exploration to fix the eigenvector closest to failing $\Sigma_\infty < D_\delta$? (Claude, AUDIT-WORKING-773921). A precise, falsifiable reader question about the matrix form, worth a sentence or a forward pointer to #deriv-causal-ib-lmi.
- **Is the “Discussion” section overloaded?** One substrate read the segment as “a chaotic mix of profound insights and dense defensive literature reviews ... shoving advanced variational bounding theorems into a Discussion segment ... severely harms readability” and would refactor the exact Lagrangian survival-imperative derivation into its own Derived segment with the active-inference review quarantined (Gemini, AUDIT-WORKING-829314). Same placement-vs-value tension flagged at #der-loop-interventional-access; preserved as a real readability signal, not a defect.

5. Candidate figures.

- **Two stacked objectives diagram.** Top layer: the scalar heuristic $\mathbb{E}[\text{value}] + \lambda \cdot \text{CIY}$, gated by identifiability; lower layer: the exact matrix form $Q_O(a) + \text{Tr}(\Lambda \mathcal{I}_O(a))$. “Drawing both makes the status drift visible — the scalar form is useful vocabulary, the exact claim belongs to the LMI/tensor version” (Codex/Claude, AUDIT-WORKING-526815, has a rendered draft).
- **U-shaped exploration-value curve.** Exploration valuation against U_M : high at high U_M (epistemic drive) and high at low U_M (survival drive), dipping in between — visualizes the two-parallel-drives composition (implicit in Claude, AUDIT-WORKING-471203).

Belongs elsewhere.

- The exact Lagrangian relaxation of the LMI and the matrix-CIY trace-product are the property of `#deriv-causal-ib-lmi` / `#deriv-causal-ib-exploration`; the Bretagnolle-Huber regret derivation is `#deriv-strategy-cost-regret-bound`. This segment is the discussion-grade surface; the theorem-grade math lives in those appendix segments.

7.6 Normative: Explicit Strategy Condition

The framework's design criterion for *when explicit strategy is worth maintaining at all*. An agent benefits from carrying an explicit strategy Σ_t (6.7) when *the cost of planning is less than the cost of learning through exploration alone*. The inequality compares the cost of constructing and evaluating the strategy plus ongoing maintenance cost (as M_t evolves and edges need revision) against the cost of learning action-outcome mappings through direct interaction plus the cost of correcting errors discovered only through execution. All costs are measured in the same units, typically time or tempo-equivalent cost.

The result is honestly *normative*, not derived. It is labeled as a design criterion rather than a theorem because loop-based and model-based approaches may differ along axes the cost inequality does not capture: final value (the model introduces bias; exploration may discover things planning cannot), risk profile (exploration risks real damage; planning risks wrong models), reversibility (some exploratory actions are irreversible), and model bias (explicit Σ_t inherits the biases of M_t ; loop-based learning does not). The inequality is correct *when* the outcomes are approximately equivalent — a condition that must be verified case by case, not assumed.

The intuitive cases are stated cleanly. The right side (explore + repair) is large — strongly favoring planning — when actions are expensive or irreversible (production deployments, military operations, surgical procedures); when exploration damage is severe (a wrong move loses the game; a wrong deployment takes down the service); when the environment is slow to respond (waiting for market or test feedback); when the action space is enormous (combinatorial planning). The left side (plan + maintain) is large — favoring pure exploration — when the environment is too complex or novel for useful models, when M_t is severely inadequate, when the environment changes faster than Σ_t can be maintained, or when actions are cheap and reversible (A/B testing, sandbox exploration).

The *normative grounding* is anchored in the persistence condition (4.5), not in an external preference postulate. The persistence condition demonstrates that agents whose correction tempo is insufficient *degrade* — descriptive, not value-laden. The cost inequality operationalizes the consequence: approaches consuming less tempo budget leave more margin above the persistence threshold. The normative element is the preference for *maintaining persistence margin*, which is hard to argue against since the alternative is degradation. The Discussion below treats the strongly-favors-planning and strongly-favors-exploration cases in detail, the persistence-margin grounding, and the connection to the three-way exploit/explore/deliberate refinement in 10.2; Working Notes flag the most useful practical reading — not as a binary decision rule but as a way to *calibrate strategy complexity* so the planning-and-maintenance cost stays just below the exploration-and-repair cost for the current environment.

Formal Expression.

An agent benefits from explicit Σ_t when:

Normative (explicit-strategy-condition).

$$C_{\text{plan}} + C_{\text{maintain}} < C_{\text{explore}} + C_{\text{repair}} \quad (7.7)$$

where: - C_{plan} : cost of constructing and evaluating the strategy (deliberation, simulation, model queries) - C_{maintain} : ongoing cost of keeping Σ_t current as M_t evolves (edge revision, structural updates) - C_{explore} : cost of learning action-outcome mappings through direct interaction (real actions, real time, real consequences) - C_{repair} : cost of correcting errors discovered only through execution (rollbacks, rework, damage)

All costs are measured in the same units (typically time or tempo-equivalent cost). The inequality requires that the two approaches produce approximately equivalent non-temporal outcomes — otherwise the comparison is between different strategies, not different approaches to the same goal.

Epistemic Status.

Normative, not derived. This is labeled *normative* because it is a design criterion (a preference for the less costly approach given equivalent outcomes), not a theorem. In practice, loop-based and model-based approaches may differ in:

- **Final value:** the model introduces bias; exploration may discover things planning cannot
- **Risk profile:** exploration risks real damage; planning risks wrong models
- **Reversibility:** some exploratory actions are irreversible
- **Model bias:** explicit Σ_t inherits the biases of M_t ; loop-based learning does not

The inequality is correct *when* the outcomes are approximately equivalent — a condition that must be verified case by case, not assumed. When the precondition fails (model-based and loop-based approaches produce qualitatively different outcomes), the cost inequality is insufficient and the choice requires richer analysis (e.g., expected regret including model error).

Discussion.

When the inequality strongly favors planning. The right side ($C_{\text{explore}} + C_{\text{repair}}$) is large when: - Actions are expensive or irreversible (production deployments, military operations, surgical procedures) - Exploration damage is severe (a wrong move in chess loses the game; a wrong deployment takes down the service) - The environment is slow to respond (waiting for market feedback, waiting for test results) - The action space is enormous (combinatorial planning problems)

In these domains, explicit Σ_t is strongly motivated even if the planning model is imperfect.

When the inequality favors pure exploration. The left side ($C_{\text{plan}} + C_{\text{maintain}}$) is large when: - The environment is too complex or novel for useful models (genuinely unknown territory) - M_t is severely inadequate (model predictions are worse than random) - The environment changes faster than Σ_t can be maintained (ρ_Σ exceeds planning capacity) - Actions are cheap and reversible (A/B testing, sandbox exploration)

Normative grounding. The cost inequality is grounded in the persistence condition (4.5), not in an external preference postulate. The persistence condition demonstrates that agents whose correction tempo is insufficient *degrade* — this is a descriptive result, not a value judgment. The cost inequality operationalizes the consequence: approaches that consume less tempo budget leave more margin above the persistence threshold. The normative element is the preference for maintaining persistence margin — which is hard to argue against, since the alternative is degradation. (In TST, the temporal optimality postulate provides an additional normative grounding specific to software development.)

Connection to the three-way tradeoff. For actuated agents, the binary explore/exploit tradeoff extends to three modes: exploit (pursue O_t via Σ_t), explore (improve M_t), and deliberate (revise Σ_t). The cost inequality addresses the coarsest question (is explicit Σ_t worth having?). The finer allocation between the three modes is addressed in 10.2 — the extended deliberation threshold is derived, but the broader three-way framing is discussion-grade. The deepest insight: deliberation is internal exploration (simulation, counterfactual reasoning, cross-domain synthesis), making it a fundamentally different resource type from external action.

Working Notes.

- The inequality as stated is static — it compares cumulative costs. A dynamic version would ask: “given current M_t quality, is it worth deliberating further or should I act now?” This connects to 4.1’s threshold: deliberation is worthwhile only when additional deliberation improves action quality enough to justify the time spent.
- Part of C_{maintain} is the cognitive cost of keeping Σ_t in the agent’s representational capacity. For LLM agents, this means fitting the strategy in the context window. A 500-node DAG may exceed this capacity, making the left side of the inequality large enough that simpler strategies (or pure exploration) become preferable despite higher exploration costs.
- The cost inequality may be most useful not as a binary decision rule but as a way to calibrate Σ_t complexity: the agent should maintain a strategy just complex enough that $C_{\text{plan}} + C_{\text{maintain}}$ stays below $C_{\text{explore}} + C_{\text{repair}}$ for the current environment. This gives a principled answer to “how detailed should my plan be?”

Incidental audit gold (lift 2026-05-30). Cross-audit “wandering thoughts” / §14-ideation harvested from the de-novo auditors’ working dirs, deduplicated across substrates and lightly attributed. Orthogonal pedagogical / framing / figure / naming material, kept separate from the certified theory-fix findings (handled elsewhere). **Coverage:** dedicated reflections at 526815, 584721, 773921, 829314, 849201 plus the batch dir 471203. Substrate attribution inferred from voice where not explicit; uncertain cases hedged.

1. Candidate Brief prose / pre-prose.

- The calibration reading (already in Working Notes) was independently singled out as “*one of the most practical pieces of advice for software engineering and organizational design I have ever seen formalized*” (Claude, AUDIT-WORKING-773921) and the segment’s most useful operational handle (Claude, AUDIT-WORKING-584721). Strong signal to promote the “*how detailed should my plan be?*” framing from Working Notes into the Brief / body.
- The reframing of the planning-vs-learning debate: the cost inequality “*reframes the question from doctrinal (planning-vs-RL camps) to empirical (which costs dominate in your domain?)*” (Claude, AUDIT-WORKING-584721). A clean contribution statement.

2. Candidate Discussion.

- **The Agile-vs-Waterfall mapping as the canonical worked instance.** Strong cross-substrate convergence (Gemini, AUDIT-WORKING-829314; Claude, AUDIT-WORKING-773921) that the four costs *are the mathematical foundation of software-engineering methodologies*: Waterfall assumes $C_{\text{repair}} \rightarrow \infty$ (space shuttle) and C_{explore} prohibitive $\rightarrow$ spend heavily on $C_{\text{plan}} + C_{\text{maintain}}$; Agile/Lean assumes C_{repair} cheap (revert a git commit) and C_{explore} fast (A/B test) $\rightarrow$ minimize planning, gather interventional data instead. “*Neither approach is universally right — the optimal strategy depends entirely on the physical parameters of the environment ... Agile in a high- C_{repair} environment like surgery dies; Waterfall in a low- C_{repair} /high-volatility web startup goes bankrupt paying C_{maintain} for a plan that is instantly obsolete.*” A candidate Discussion example or TST cross-reference.
- **The LLM context-window as a structural C_{maintain} amplifier.** The Working-Notes point (a 500-node DAG may exceed an LLM’s representational capacity) was read as load-bearing: “*for an LLM, C_{maintain} is massive because of the finite context window — every plan step consumes tokens, crowding out operational memory ... the math says an LLM should not write a master plan up front; it should just start writing code and see what breaks — pure exploration is mathematically optimal for them given their current cognitive costs.*” And the engineering corrective: “*if we want LLMs to plan better, we must artificially reduce their C_{maintain} — external structured scratchpad / hierarchical memory to offload the description length of the DAG; AAT gives the exact equation for why this intervention is necessary*” (Gemini, AUDIT-WORKING-829314). Candidate Discussion sharpening connecting the cost inequality to scaffolding design — and it converges with the persistence-bandwidth-floor framing already noted in #impl-causal-access.

3. Follow-up items (soft / scope-honesty).

- **“Loop-based learning does not inherit biases” overstates.** The Epistemic Status contrasts explicit- Σ_t (inherits M_t bias) with loop-based learning (does not). But direct exploration is still biased by policy selection, partial observability, confounding, update rules, and the model used to *interpret* observations. The honest contrast is “*planning inherits model bias differently*,” not “*loop learning is unbiased*” (Codex/Claude, AUDIT-WORKING-526815). Candidate one-clause fix.
- **Persistence-margin grounding is not universal across continuity stances.** The normative grounding leans on “preference for persistence margin, hard to argue against.” But the continuity-stance work (#disc-continuity-stance) recognizes task-terminal / indifferent / negotiated-continuation agents — the criterion is compelling for agents whose *objectives require ongoing persistence*, and should not silently override the objective’s stance toward continuation (Codex/Claude, AUDIT-WORKING-526815). Candidate scope qualifier.
- **Dependency-declaration gap.** The criterion is grounded in #result-persistence-condition, which is not declared in depends:. Acceptable if Discussion-only dependencies aren’t tracked, but the grounding is substantive (Codex/Claude, AUDIT-WORKING-526815). Worth a glance against the project’s depends-tracking policy.
- **Common-units / risk treatment.** Irreversible damage and heavy-tailed repair costs are not well represented by a static scalar inequality unless the cost measure already embeds risk; the equivalent-non-temporal-outcomes precondition is doing a lot of work (Codex/Claude, AUDIT-WORKING-526815). The body’s caveat already names this; the signal is that it is easy to read past.

4. Readers often ask / wonder.

- **How does the three-way budget (exploit / explore / deliberate) get allocated in real time?** Several substrates noted the body’s binary inequality answers only the coarsest question and the real interest is the dynamic three-way allocation — “*this maps exactly to the control loop of an advanced autonomous agent (Devin / SWE-agent); understanding how the framework triggers transitions between the three modes is the key to building AAT-compliant agents*” (Claude, AUDIT-WORKING-773921; Gemini, AUDIT-WORKING-829314 — classical RL treats deliberation as computationally free, AAT costs it via C_{plan}). The body forward-points to #disc-exploit-explore-deliberate; readers want that pointer prominent.

Belongs elsewhere.

- The compounding-fragility-of-long-chains material (chain-confidence decay, evidence starvation, cognitive cost = a “triple depth penalty”) that auditors raised while reading this segment belongs at #der-chain-confidence-decay (the next segment), not here.

7.7 Derivation: Mechanism Counterfactuals Separate Strictly From Level 3 — Witness, Reducibility Boundary, and Internalization

Two SCMs can agree on *every* Level-1/2/3 quantity — the identical joint law of all hard-intervention counterfactual worlds — while assigning probabilities 1 and 0 to the same noise-preserving mechanism-replacement query; the separation is therefore strict in exactly the Causal Hierarchy Theorem’s sense, and it is generic rather than exotic: every SCM whose exogenous space admits a measure-automorphism with realized effect on some mechanism disagrees with *its own relabeling* on a latent-anchored query (Result 4). The boundary is characterized by the *specification language* of the imagined mechanism — replacements written over potential-outcome coordinates (the closure of the Level-3-exposed channels) reduce to Level 3, while replacements anchored to the latent parameterization need not, and the witness’s provably does not — and internalization via selector variables re-represents but provably does not collapse the separation.

Formal Expression.

Throughout, $M = \langle U, V, F, P(U) \rangle$ is a finite recursive (acyclic) SCM. The **Level-3 content** of M is the joint law of the potential-outcome process $\text{PO}(M) = \{V_w(\cdot) : w \text{ a hard intervention on a subset of } V\}$ on $(U, P(U))$ — by [1, Defs.~7, 9] this is precisely the family of $\mathcal{L}_3$ valuations $P^M(\mathbf{y}_x, \dots, \mathbf{z}_w)$, and it subsumes all $\mathcal{L}_1, \mathcal{L}_2$ valuations and nested counterfactuals. Write $M \equiv_3 M'$ for equality of Level-3 content.

[Definition (noise-preserving mechanism replacement; after Pearl’s local surgeries — [2, §7.2.4]: actions “represented as local surgeries,” respecifying “those few mechanisms that are perturbed” with “all other mechanisms... unaltered”; the atomic case defined by mechanism replacement in [3, §2])]]

For $g : \text{dom}(pa_i) \times \text{dom}(U) \rightarrow \text{dom}(v_i)$, let $M[f_i := g]$ denote M with the i -th structural equation replaced and $\langle U, P(U) \rangle$ unchanged. **Mechanism-counterfactual queries** are probabilities of Boolean combinations of events over worlds drawn from the submodels of both M and $M[f_i := g]$, all sharing the exogenous draw u ; $\text{do}(X = x)$ is the constant special case, and evaluation extends Pearl’s abduction-action-prediction three-step verbatim.

There exist $M_A \equiv_3 M_B$ over a shared signature $\langle U, V, P(U) \rangle$ and a mechanism-counterfactual query q with $q(M_A) = 1$, $q(M_B) = 0$. Hence noise-preserving mechanism counterfactuals are not Level-3 quantities.

Derived (Result 1 — strictness witness).

Derivation. Binary U_X, U_Y independent uniform; $X = U_X$ in both models;

$$M_A : Y = X \oplus U_Y, \quad M_B : Y = X \oplus U_Y \oplus 1. \quad (7.8)$$

The response pair $(Y_{x=0}, Y_{x=1})$ is $(U_Y, 1 \oplus U_Y)$ in M_A and $(1 \oplus U_Y, U_Y)$ in M_B ; the $P(U)$ -preserving bijection $\varphi(u_X, u_Y) = (u_X, 1 \oplus u_Y)$ carries one potential-outcome process onto the other, so the joint law of $\text{PO}(\cdot) - (Y_{x=0}, Y_{x=1})$ uniform on $\{(0, 1), (1, 0)\}$, independent of $X = U_X$, with the actual $Y = Y_X$ and the trivial responses under interventions on Y — is identical: $M_A \equiv_3 M_B$. Now replace Y ’s mechanism by $g(x, u_Y) = u_Y$ and take $q = P(Y_{f_Y:=g} = Y_{x=0})$. In M_A , $Y_{f_Y:=g} = U_Y = Y_{x=0}$ a.s., so $q = 1$; in M_B , $Y_{f_Y:=g} = U_Y = 1 \oplus Y_{x=0}$ a.s., so $q = 0$. The evidence-conditioned form behaves identically: $P(Y_{f_Y:=g} = 1 \mid X = 0, Y = 1)$ is 1 in M_A and 0 in M_B . ■

Call a replacement **PO-specified** when the new mechanism is $v_i := h(pa_i, \pi(\cdot), \varepsilon)$ for fixed measurable h , fresh noise $\varepsilon \perp U$ of given law, and $\pi(u)$ a fixed finite tuple of potential-outcome coordinates — each entry $a(v_j)_w(u)$ for specification-fixed variable j and hard intervention w — the closure of the Level-3-exposed channels. This includes as special cases every **L3-specified** replacement $v_i := h(pa_i, f_i(pa_i, \cdot), \varepsilon)$: deterministic parent-only re-wirings, independent randomized devices, all shift-scale soft interventions $a(pa_i)f_i + b(pa_i)$, and any replacement reading named counterfactual quantities of the original model (e.g. “output $Y_{x=0}$ ’s value”). Every cross-world query over M and $M[f_i := h]$ with h PO-specified is determined by the Level-3 content of M together with the law of ε ; PO-specified replacements therefore never separate a $\equiv_3$ pair.

Derived (Result 2 — reducibility of PO-specified replacements).

Derivation. By topological induction on the modified world’s variables: upstream variables retain their original potential responses; $v_i^h = h(pa_i^h, \pi(u), \varepsilon)$, where each coordinate of $\pi(u)$ is by definition an element of $\text{PO}(M)$ and $f_i(pa_i^h, u)$ — where read — is the original potential outcome of v_i under the hard intervention setting its parents to the already-computed values, an element of $\text{PO}(M)$ at a computed index; downstream variables are $V_{w \cup \{v_i=v_i^h\}}(u)$, again elements of $\text{PO}(M)$ at computed indices. The modified-world trajectory is thus a measurable functional of $\text{PO}(M)$ and ε , so every joint law with original-model worlds is a functional of Level-3 content. Model-invariance is by construction: a PO-specification names *counterfactual coordinates*, which $\equiv_3$ -equivalent models value identically in law — unlike a latent-written table, which names points of $\text{dom}(U)$ and acquires a PO-reading only model-relatively. ■

For any finite family $\mathcal{G} = \{g_1, \dots, g_k\}$ of replacements for f_i , the enlarged SCM $M^\dagger$ with selector $S \in \{0, \dots, k\}$ (endogenous, with constant mechanism $S := 0$, so that $Y_{S=j}$ is a hard-intervention quantity) and $f_i^\dagger(pa_i, s, u) = f_i(pa_i, u)$ if $s = 0$, else $g_s(pa_i, u)$ — Pearl’s interventions-as-variables construction ([2, §3.2.2, Eq.~(3.7), Fig.~3.2]; originally [3, Eq.~(3), Fig.~1]), whose $x_i = I(pa_i, f_i, u_i)$ with $I(a, b, c) = f_i(a, c)$ whenever $b = f_i$ is precisely mechanism replacement internalized as an added parent variable — “applicable to any change in the functional relationship f_i and not merely to the replacement of f_i by a constant” (§3.2.2) — applied reflexively — satisfies $Y_{f_i:=g_j} = Y_{S=j}$: every mechanism counterfactual over $(M, \mathcal{G})$ is a hard-intervention (Level-2/3) quantity of $M^\dagger$. But the map $M \mapsto M^\dagger$ is not a functional of Level-3 content: by Result 1, $M_A \equiv_3 M_B$ while $M_A^\dagger \not\equiv_3 M_B^\dagger$ (their Level-3 contents already disagree on $P(Y_{S=1} = Y_{X=0})$). Internalization is representation change over an enlarged variable set, exactly as Level 3 is “Level 1 over the potential-outcome event space” — the informational hierarchy, defined relative to a fixed variable set, survives. ■

Derived (Result 3 — internalization re-represents, does not collapse).

Let ϕ be a $P(U)$ -preserving bijection of $\text{dom}(U)$ and let M^ϕ be the **relabeling** of M : same signature, mechanisms $f_j^\phi(pa_j, u) := f_j(pa_j, \phi(u))$ for every j . Then (i) $M \equiv_3 M^\phi$ always; and (ii) taking the latent-written replacement table $g(pa_i, u) := f_i(pa_i, \phi(u))$ and the query $q = P(V_i^{[f_i:=g]} = V_i)$,

Derived (Result 4 — genericity: every anchor-mobile SCM separates from its own relabeling).

$$q(M^\phi) = 1, \quad q(M) = 1 - \mu_\phi^i, \quad \mu_\phi^i := P\{u : f_i(pa_i(u), \phi(u)) \neq f_i(pa_i(u), u)\}, \quad (7.9)$$

so any (ϕ, i) with realized disagreement mass $\mu_\phi^i > 0$ yields an $\equiv_3$ pair separated by a latent-anchored query. Call M **anchor-mobile** when such (ϕ, i) exists and **anchor-rigid** otherwise. Every SCM with an exogenous coordinate uniform (more generally, exchangeable) over two or more values on which some mechanism genuinely depends at realized parent values is anchor-mobile via a transposition; the witness pair is the minimal instance $(\phi : u_Y \mapsto 1 \oplus u_Y)$. Anchor-rigidity is thus confined to SCMs whose noise is effectively inert or whose atom masses are fully asymmetric (admitting no nontrivial automorphism; for nonatomic noise, automorphisms always exist, so rigidity there reduces to inertness alone) — the separation is generic, not exotic.

Derivation. (i) By induction along the topological order, for every hard intervention w the solution map of M^ϕ ’s w -submodel at u equals that of M ’s at $\phi(u)$ (unreplaced mechanisms of M^ϕ evaluate M ’s at $\phi(u)$; intervened equations are the same constants in both), so the PO-tuple map satisfies $\text{PO}(M^\phi)(u) = \text{PO}(M)(\phi(u))$ pointwise, and the pushforward laws coincide because ϕ preserves $P(U)$. (ii) In M^ϕ the table g is f_i^ϕ , so the replaced world is the actual world and $q = 1$. In M , acyclicity keeps V_i ’s parents at their actual values in the replaced world, so $V_i^{[f_i:=g]}(u) = f_i(pa_i(u), \phi(u))$ against $V_i(u) = f_i(pa_i(u), u)$, giving $q(M) = 1 - \mu_\phi^i$. ■

Corollary (anchoring = choice of coordinates). M^ϕ is measure-isomorphic to M — the separated pair differ in *nothing but the parameterization of the latent space*. Latent-anchored mechanism-counterfactual content is therefore exactly commitment to coordinates on U : real content (the queries have definite, disagreeing values), invisible to all three layers, and supplied only by the committed representation.

In M_A , the separating $g(x, u_Y) = u_Y$ coincides with the PO-coordinate $Y_{x=0}$ — a PO-specification — while relative to M_B the same latent-written table reads as $Y_{x=0} \oplus 1$: a latent-written table determines a PO-specification only *relative to a model*, and $\equiv_3$ -equivalent models translate it inconsistently. The strictly-beyond-Level-3 class is therefore not “formulas mentioning u ” but replacements whose meaning is anchored to the latent parameterization itself — *hold this very background fixed and change the law over it*. Result 2 gives the exact sufficient class for reducibility (uniform PO-specifications); Result 4 shows the non-reducible remainder is generically inhabited. What remains open is the converse characterization: whether every replacement whose queries are $\equiv_3$ -invariant admits a uniform PO-specification (see Discussion).

Discussion (the boundary, and why Results 1, 2, and 4 are consistent).

Epistemic Status.

Exact. Results 1-4 are elementary, self-contained derivations under the stated assumptions (finite recursive SCMs; Result 2 additionally assumes the replacement is PO-specified). The Level-3 semantics is imported and was primary-verified against the on-disk text of [1, Defs.~2, 5, 7, 9 and Thm.~1] (ref/bareinboim-2022-pearl-hierarchy.pdf, 2026-07-16). The surgery and interventions-as-variables machinery is Pearl's, primary-verified 2026-07-16 against both the 1994 primary ([3, §2], ref/pearl_calculus_of_actions_1994.pdf: atomic interventions defined by mechanism replacement; Eq.~(3)/Fig.~1 internalization) and the second-edition text ([2, §§3.2.2, 7.2.4], read directly: Eq.~(3.7)/Fig.~3.2 with the "any change in the functional relationship" scope note; §7.2.4's local-surgery semantics — one mechanism perturbed, locality in mechanism space). Noise-dependent replacement mechanisms are included in Pearl's own construction, so the internalization used in Result 3 is literally his move and the new content is its *non-collapse* reading. What is new here (per the math-novelty discipline, stated without deflation): the strictness witness over full- $\equiv_3$ pairs, the PO-closure reducibility boundary, the genericity-via-relabeling theorem with its anchoring-as-coordinates corollary, and the non-collapse reading of internalization. Priority: a professional literature search (Undermind, 2026-07-16, estimated full coverage of the relevant population; report retained as a working aid, see Working Notes) found **no published statement or implication of the separation**, in either the two-model-witness or the exogenous-relabeling form; the nine nearest papers were then read in full against their primary text (2026-07-16) — the danger scan for a buried separation lemma (through Ibeling and Icard 2020/2023, Shpitser and Pearl 2008, Cabreros and Storey 2019, de Lara 2023, and Aliprantis 2015 as well as the three nearest below) came back empty, and the three nearest results *position* rather than anticipate it. Bongers, Peters, Schölkopf and Mooij (2016/2021, §4) introduce **exogenous reparametrization** as an operation on SCMs and prove it preserves *observational and interventional* semantics (their "causal semantics"; the general operation of which our measure-preserving ϕ is the automorphism special case, and Prop. 53's "symmetries acting on the space of SCMs" is the group-action framing Result 4 uses) — they do not treat the counterfactual layer, so the Level-3 preservation under relabeling is *our* Result 4(i) and its failure at latent-anchored mechanism counterfactuals is Result 4(ii): the operation is theirs, the layer-3 boundary of its semantics-preservation is new here. Chen and Du (2025) define **exogenous isomorphism** ($\sim_{EI}$; Def. 3.1 — a *component-wise* latent bijection $h = (h_i)$ with $P_U^{(2)} = h_{\#}P_U^{(1)}$ and mechanism isomorphism $f_i^{(2)}(v, h_i(u)) = f_i^{(1)}(v, u)$) and prove $\sim_{EI} \Rightarrow \sim_{\mathcal{L}_3}$ (Thm. 3.2), positioning $\sim_{EI}$ -identifiability as strictly stronger than complete counterfactual identifiability. The witness pair — and every component-wise relabeling of Result 4 — is exogenously isomorphic in this *exact* sense (take $h_Y = \oplus 1$, $h_X = \text{id}$), which is *stronger* than the measure-isomorphism the corollary asserts in general: so even their strongest identification target, $\sim_{EI}$, leaves latent-anchored mechanism-counterfactual answers open — the obstruction lives in the specification language, not the model class. De Lara (2025, Def. 17 / Thm. 28) characterizes the counterfactual-equivalence ($\equiv_3$) class of an SCM and disentangles its "arbitrary and unfalsifiable" counterfactual coupling from the testable interventional constraints — the same conventionality structure this segment exhibits one rung higher, over mechanism-change queries on the $\equiv_3$ -complete class; his footnote 9 explicitly routes through Chen-Du's $\sim_{EI}$, so the three nearest works form one connected cluster, none of which states the separation. Older adjacencies as previously recorded, now primary-confirmed: the underdetermination of F by counterfactuals is folklore (Aliprantis 2015's "many DGPs, same potential outcomes"; canonical response-function representations, Balke and Pearl 1994) but exhibits no query the models disagree on; Saha, Rathore and Garain 2025 claim soft interventions are "strictly more expressive," a query-language expressiveness claim whose shift-scale class is PO-specified, hence Level-3-reducible by Result 2. Result 4 replaces Result 1's bare existence with a genericity statement in the CHT's register: separation holds for every anchor-mobile SCM, with anchor-rigidity confined to inert-noise or automorphism-free-noise cases; whether the anchor-rigid remainder also separates by other constructions is subsumed in the open converse below.

Discussion.

What the separation means. Level 3 fixes the joint law of all hard-intervention worlds; it does not fix how those worlds are co-registered against *alternative laws on the same background*. Both witness models even agree on the post-replacement interventional distribution ($Y_{f_Y:=g} \sim \text{Bern}(\frac{1}{2})$ in each); the disagreement lives exclusively in the cross-world coupling between actual-law and replaced-law worlds through the shared u — content invisible to every experiment and every standard counterfactual. This licenses the hierarchy extension in precisely the CHT’s relative sense: “Level 4” quantities (latent-anchored mechanism counterfactuals) are strictly more informative than Level 3 relative to the model’s variable set, and — like every level — re-representable as lower-level over an enlarged set without that fact collapsing anything (Result 3). Result 4’s corollary sharpens what the extra content *is*: since the separated pair are relabelings of one another, latent-anchored content is exactly a commitment to coordinates on the latent space — nothing more, and (because the queries then have definite values) nothing less.

The open converse. Result 2 shows uniform PO-specifications are $\equiv_3$ -invariant; the converse — whether every replacement whose queries are invariant across the $\equiv_3$ -class admits a uniform PO-specification — is open and is the sharp remaining question about the boundary. A plausible route is an orbit-averaging argument over the relabeling group of Result 4, but invariance of a query’s *value* does not obviously yield a *pointwise* uniform representation of the specification; a counterexample would itself be informative (an invariant-but-unrepresentable replacement would name a fourth-layer quantity computable from Level 3 in value yet not in form). One exhaustive datapoint (independent verification, 2026-07-16): on the witness signature’s minimal universe (all 4096 SCMs; the 24-model $\equiv_3$ -class of M_A ; all 256 replacement tables), every $\equiv_3$ -invariant replacement — there are exactly 4 — is uniformly PO-representable, so the converse holds there; the value-vs-pointwise-form obstruction remains untouched in general.

Consequence for structural imagination. The boundary splits imagination-as-mechanism-change into *navigation* (L3-specified changes — re-wirings, new independent devices, output-shifting policies — computable from counterfactual-grade knowledge) and *positing* (latent-anchored changes — “this very situation under a different law” — whose answers no experimentable or counterfactual content determines, and which are well-defined only relative to a committed representation F among the $\equiv_3$ -class). 7.8 develops the agent-architecture consequences: the imagination workspace’s grounding tag marks, for latent-anchored content, information that observation cannot even in principle supply.

Relation to the do-operator’s special-case status. $do(X = x)$ is the constant replacement — both parent-only and latent-free — so all of Pearl’s hierarchy sits on the reducible side of Result 2’s boundary, which is why the question “is mechanism change a fourth rung?” could not be settled by inspecting the standard machinery: the standard machinery never leaves the L3-specified class.

Working Notes.

- **Provenance.** Landed 2026-07-16 from spikes/spike-causal-level-4-strictness-2026-07-16.md (commissioned same day; owed-items discharge — primary-source verification, Result 2 generalization, Result 4 — same day, recorded in the spike’s addendum), which resolves the strictness question left open by spikes/.integrated/spike-causal-level-4-formal.md (2026-03-14, Approach 1: the $do_s(f_X := g)$ operator and well-definedness, without a witness).
- **Naming caution.** Prose should say “mechanism counterfactuals separate strictly from Level 3” or “latent-anchored mechanism counterfactuals”; reserve bare “Level 4” for contexts where the relative-to-variable-set sense is already established, to avoid reading as a claim that Pearl’s hierarchy was somehow incomplete on its own terms.
- **Priority-search record.** The Undermind report (9 papers, clean miss) is at ref/layer-3-mechanism-counterfactual-sep.md — a working aid, not segment-citable; all nine are in relata with PDFs. The three nearest priors (Chen and Du 2025; Bongers et al. 2016/2021; de Lara 2025) and the six adjacencies were read in full against primary text 2026-07-16 (relata verify events recorded), discharging the ES characterizations from abstract-grade to primary-confirmed — including the one correction that mattered: Bongers prove interventional, not counterfactual, preservation, so the Level-3 preservation is Result 4(i),

not theirs. Standalone-note decision is Joseph's (ops queue); the practicals live in #disc-law-discovery-ceiling Discussion, written to lift.

7.8 Discussion: Structural Imagination — Model-Modifying Reasoning and the Scope Boundary of Directed Separation

Pearl's three levels all reason *within* a fixed structural causal model; a class of reasoning AAT's agents demonstrably perform — imagining alternative causal structures, planning actions that *change* the environment's mechanisms, and evaluating aspirational targets — operates *on the space of models* instead. This segment names that task class, locates its inferential status via the strict separation derived in 7.7 (latent-anchored mechanism counterfactuals exceed Level 3; L3-specified ones do not), locates the scope boundary of directed separation exactly at its edge — goal-blind processing is correct for within-model evidence, and goal-*directed* exploration is correct (not a failure) for model-space reasoning — and proposes the tagged-workspace structure $X_t = (M_t, G_t, W_t)$ under which the two modes coexist without contamination. Hallucination is then diagnosed as a *level-confusion* failure: an explored model promoted to M_t without grounding — a mechanism complementary to the quantitative goal-coupling displacement bound of AB.

Formal Expression.

Pearl's SCM machinery already represents mechanism replacement ([2, §7.2.4] “local surgeries”): for $M = \langle U, V, F \rangle$, define $M[f_X := g]$ by replacing one structural equation, with $do(X = x)$ the constant-function special case. An action a is **structure-modifying** with respect to the environment's causal model M when the post-action environment is described by $M' = \sigma(M, a) \neq M$ — the structural equations themselves have changed (a bridge built, a law enacted, a program written), as opposed to standard actions, which set variable values within fixed mechanisms. The corresponding planning query optimizes over model space:

Formulation (structure-modifying action; sketch).

$$a^* = \arg \max_a \mathbb{E}_{M_{t+1} \sim P(\cdot | a, M_t)} [J(M_{t+1})], \quad (7.10)$$

a task with no Level-1–3 analog: the search space is a mechanism space, not a value space, and evaluation marginalizes over the models an action could produce.

The task class is inferentially two-sided, per the exact results of 7.7. Mechanism changes *specified over Level-3-exposed channels* (re-wirings of parents, independent new devices, output-shifting policies) are Level-3-computable: an agent with counterfactual-grade knowledge can evaluate them without further commitment. Mechanism changes *anchored to the latent background* — “this very situation under a different law” — separate strictly from Level 3: two models agreeing on every associational, interventional, and counterfactual quantity can disagree maximally on them, so their answers are underdetermined by all experimentable content and exist only relative to a committed structural representation. On top of that informational distinction, the task class remains *computationally and cognitively* distinct (mechanism-space search, no analog of the do-calculus for identification, inseparability of model identification from inference when actions modify the model), and it is precisely the reasoning mode AAT's structural-adaptation machinery (4.8, 2.4) presupposes an agent can perform.

Discussion — inferential status; separation derived at #deriv-mechanism-counterfactual-separation.

The complete agent state extends to $X_t = (M_t, G_t, W_t)$, where W_t is the **imagination workspace**: the set of currently-explored alternative models (candidate structures, aspirational futures, hypothetical mechanism changes), explicitly tagged as ungrounded. Directed separation (6.3) holds for M_t by definition of what M_t is — evidence-processing remains goal-blind — and *correctly fails* for W_t : which alternative structures to explore is inherently goal-directed, and

that coupling is the productive core of planning, construction, and aspiration rather than a leak to be bounded. Falsification handle: an architecture that maintains the tag (grounded vs explored) should show the bias-bound pathologies of goal coupling only in its W_t -to- M_t promotion step, not in its M_t updates.

Hypothesis (tagged workspace).

Maintaining an infeasible aspirational objective can dominate feasible targets by the feasible objective's own measure — see the aspiration hypothesis landed in 8.4's Discussion, which modifies the reading of "objective revision is the last resort."

Hypothesis (aspiration; formal statement at #def-satisfaction-gap).

Epistemic Status.

Discussion-grade throughout; nothing here is derived. The structure-modifying action formulation is a sketch on Pearl's own machinery (imported: Pearl 2009 §7.2.4; hierarchy strictness Bareinboim et al. 2022 — both already recapitulated in 7.1); the planning query is a formulation; the W_t workspace and the level-confusion diagnosis of hallucination are hypotheses with a stated falsification handle; the hierarchy status of the task class is settled and *exact* at 7.7 (strict separation for latent-anchored mechanism counterfactuals; Level-3 reducibility for L3-specified ones; internalization as re-representation, not collapse). The typing of what $\sigma(M, a)$ acts on interacts directly with the epistemic-target-ontology decision (PROPOSALS SP-30: $S_t = (\Omega_t, \theta)$ gives law-content its own slot, and a structure-modifying action is then an action on world-side θ) — the vocabulary here is stated in current canon terms and should be re-typed when SP-30 is decided. Adjacent external literature (causal abstraction: Beckers & Halpern; Rubenstein et al.; transportability: Bareinboim & Pearl) is flagged for a relata pass before this segment advances past draft.

Discussion.

Where directed separation's scope boundary actually sits. 6.3 already distinguishes goal-directed event *selection* (legitimate) from goal-blind event *processing* (the invariant). This segment adds a third distinction on the *content* axis: processing that updates the grounded model M_t (goal-blind, Levels 1–3) versus reasoning that explores ungrounded alternatives W_t (goal-directed, structure-modifying). The κ problem for merged architectures then reframes: the issue is not the amount of goal-coupling but **cross-contamination between two modes that individually have correct, opposite coupling disciplines** — an LLM's context window holds grounded belief and explored hypothesis in the same representation with no tag, so Level-4-appropriate coupling bleeds into Level-1–3 processing. Keeping the modes distinct is a design problem, not a parameter-tuning problem.

Hallucination as level confusion — complementary to the bias bound. Canon's quantitative hallucination treatment (AB and its NeurIPS Paper 3 sharpening) bounds the displacement of M_t updates under goal coupling — a *drift* mechanism. The level-confusion diagnosis names a different failure: a W_t element (imagined structure, hypothetical world) promoted to M_t status without grounding — a *mis-typing* mechanism. The two are complementary: drift corrupts the grounded model gradually through coupled processing; promotion corrupts it discretely through a missing tag. The light side of the same machinery is imagination — the tag maintained, exploration goal-directed, promotion gated on evidence.

Relation to structural adaptation. 4.8 establishes when parametric updates within the current model class must give way to structural change, treating the event discretely. Structural imagination is the *capability* that responding to that event presupposes: generating and evaluating candidate structures is model-space reasoning. A continuous structural-adequacy mismatch ($\delta_{\text{structural}}$, sitting between epistemic update and feasibility evaluation in the orient cascade) was sketched in the originating spike and is deliberately *not* landed here — it needs its own derivation of measurability before it earns a mismatch-family slot.

Working Notes.

- **Provenance.** Landed 2026-07-16 from spikes/.integrated/spike-causal-level-4.md + spike-causal-level-4-formal.md (2026-03-13/14, Joseph's aphorism-prompted overnight cluster) via the bulk-64 verification queue; the strictness question those spikes left open was resolved the same day by spikes/spike-causal-level-4-strictness-2026-07-16.md, landed as 7.7.
- **SP-30 interaction (re-type on decision).** If $S_t = (\Omega_t, \theta)$ is adopted, restate: a structure-modifying action is an intervention on world-side θ ; W_t explores candidate θ -content; the grounded/explored distinction becomes a typing on the agent's θ -beliefs. If state-only Ω is kept, the current phrasing stands.
- **Relata pass owed** before promotion: Beckers & Halpern (causal abstraction), Rubenstein et al., Bareinboim & Pearl transportability, and the mental-simulation/prospection literature the spike gestured at — none load-bearing for the discussion-grade claims, all needed for honest Related-Work grounding.
- $\delta_{\text{structural}}$ (fifth mismatch type) deliberately withheld — see Discussion; a future spike would need to derive what makes structural inadequacy *measurable* as a continuous signal before the cascade gains a slot.

7.9 Discussion: The Law-Discovery Ceiling — Conventionality, Positing, and Exogenous Intervention for Embedded Agents

The strict separation of 7.7, read from the agent's side, becomes a statement about what embedded agents can and cannot learn of their world's laws: counterfactually-complete inquiry converges to a Level-3 equivalence class of models, never to the law itself, and the surplus — the structural functions and the coordinatization of the exogenous background — is generically underdetermined by all experimentable content. This segment develops the agentic consequences: the ceiling as a candidate instance of the identifiability-floor pattern at the top of the causal hierarchy; the *conventionality license* (within the equivalence class, commitment to a provisional law is rationally free for all navigation purposes, so the honest selection criteria are compression, tempo, and communicability); the *positing regress* (achievable mechanism-change is navigation one level down, so irreducible positing coincides with the boundary of empiricism); a three-kind taxonomy of exogenous intervention with distinct in-world visibility signatures; and the *testimony channel* — the one remaining channel type by which an inside agent could hold true rather than merely committed positing-content. Throughout, claims marked *derived* inherit exactness from 7.7; everything else is discussion-grade reading, and the whole segment carries the standing caveat that it operates inside the structural-causal-model frame — reading a natural environment as an SCM is itself a modeling commitment, so every claim below is conditional on that frame applying.

*Formal Expression.***The ceiling.**

Let an agent's inquiry converge, in the limit of unbounded observation, intervention, and counterfactual inference, to the complete Level-3 characterization of its environment. By the separation theorem this limit determines only the $\equiv_3$ -equivalence class: models agreeing on every associational, interventional, and counterfactual quantity may still differ on latent-anchored mechanism-counterfactual queries, and by the genericity result (Result 4 there) essentially every environment model with genuine stochasticity differs from its own exogenous relabeling in exactly this way. **World-law discovery via the causal hierarchy therefore has a proven ceiling:** its asymptote is an equivalence class, and membership in the class exhausts what any experiment-driven process can establish. The residue — which structural decomposition, which coordinatization of the background — is not a further fact awaiting better instruments; it is unreachable in principle by the agent's channel types.

Derived (at #deriv-mechanism-counterfactual-separation); agentic restatement.

In AAT's cross-sectional vocabulary this is a candidate instance of the identifiability-floor pattern (5.1) sitting at the top of the causal hierarchy rather than inside any estimation problem: information (mechanism identity beyond $\equiv_3$) that no channel available to the embedded agent can recover, with the characteristic floor structure of an exactly-described unrecoverable residue and named escapes. The escapes here are structural rather than statistical: authored access (the agent holds the coordinatization because it built the system — see AW) and testimony (below). Promotion to the M1 instance catalog would require the full F-catalog treatment and is deliberately not claimed here.

Discussion — candidate floor instance.

The conventionality license.

Every query the agent needs for acting in the world as it stands — prediction, intervention planning, policy evaluation, regret — is Level-3 content, invariant across the class. Hence two agents committed to different provisional laws with identical $\mathcal{L}_3$ -content act identically on every realizable task: **within the equivalence class, commitment is fitness-free for navigation**. Two consequences follow. First, the honest selection criteria for a provisional law are engineering criteria — description length, inferential tempo, communicability — because truth does not adjudicate within the class; parsimony here is not a metaphysical wager but the correct optimization of a genuinely free choice. Second, the class boundary supplies a **disagreement-hygiene test**: a dispute between two committed models is *convictable* when their $\mathcal{L}_3$ -contents differ (an experiment can pay off) and *conventional* when they coincide (no experiment can ever pay off, and the honest resolution is an explicit coordinate choice). The PO-closure result gives the working form of the test: if the disputed content is expressible as a uniform PO-specification, it is on the convictable side.

Derived consequence + Discussion.

The posing regress.

Distinguish, per 7.8, *navigation* (mechanism changes specified over Level-3-exposed channels — evaluable from counterfactual-grade knowledge) from *posing* (changes anchored to the actual latent background — representation-committed). For an embedded agent, every **achievable** mechanism change is navigation one level down: damming a river, rewiring a circuit, amending a constitution are interventions in a finer-grained model in which the coarser model's "law" is configuration. Mechanism-change at level k is value-setting at level $k + 1$; the internalization result is this regress's formal engine. The regress terminates only at the environment's ultimate laws, where no finer level is available to the agent — and that terminus is exactly where the ceiling says experimentable content gives out. **The boundary of irreducible posing coincides with the boundary of empiricism**. One corollary worth stating plainly: whether an environment's randomness is *true* or *fated* (a deterministic function of hidden coordinates, as in an authored world's keyed noise) is itself a posing-class question — any fated coordinatization reproducing the same $\mathcal{L}_3$ -content is a relabeling in the theorem's sense — which is why that question remains permanently open from inside. (The scope is exact $\equiv_3$ -equivalence: accounts that alter testable content, as local hidden-variable theories do under Bell-type constraints, are convictable in the ordinary way and outside the class.)

Formulation.

Exogenous intervention — three kinds, three visibility signatures.

Suppose the environment admits an author — an agent holding the latent coordinatization, of which the authored-world case (AW) is the constructive existence proof. Author action decomposes into three kinds with strictly ordered in-world visibility:

Formulation.

Table 7.3

<i>Kind</i>	<i>Author-side operation</i>	<i>In-world visibility signature</i>
Selection	Choose among law-consistent worlds (rejection sampling against a predicate)	None, in principle: every world in the support is fully law-consistent; no inquiry at any hierarchy level distinguishes a selected world from an unselected one
Value-setting intervention	Set variable values mid-history	<i>Registerable</i> as anomaly (broken learned statistics, violated conservation checks) but <i>unidentifiable in kind</i> : any intervention history is re-representable as ordinary law of an enlarged model, and the enlargement is not a functional of $\mathcal{L}_3$ -content — the intervention-reading and the bigger-law-reading are coordinatizations of the same experimentable content
Mechanism replacement (latent-anchored)	Change a law mid-history holding the actual background fixed	Effects may be arbitrarily visible; the event's <i>identity</i> — same background, new law, versus different background, versus larger law throughout — is latent-anchored content the inside agent provably cannot resolve

The middle row is a sharpening of the classical epistemology of miracles from an evidentiary difficulty into an identifiability statement: it is not that testimony for an intervention is weak evidence, but that even counterfactually-complete evidence cannot force the intervention-reading over the enlarged-law-reading. An honest agent in a possibly-authored environment therefore carries a model with an **intervention-selector slot whose semantics it cannot pin** — natural variable, author-channel, or absent — and holding that slot open is the exact truth of its situation rather than indecision. The commitment it makes about the slot is free relative to all possible data yet consequential for its positing-grade reasoning.

Discussion.

The testimony channel.

The ceiling characterizes a content class that exists (the author-side surplus), matters (it answers the positing query class, including the identity of Kind-2/3 events and the fated-randomness question), and is unreachable by every empirical channel. If an inside agent is to hold *true* rather than merely *committed* positing-content, exactly one channel type remains: **testimony from a holder of the coordinatization**. Two structural corollaries: (i) verification of testified positing-content cannot be empirical — by construction — so it is either relational (trust in the source) or *transferred*, when the testimony bundles convictable $\mathcal{L}_3$ -content whose success lends credibility to the unconvictable remainder; (ii) *communication through the anomaly stream is representation-independent* — an intervention pattern can carry a message whose content is Layer-1-observable, so reception and decoding work without settling what the channel is; only the interpretation of the channel's nature requires the commitment.

Discussion.

Epistemic Status.

Discussion-grade throughout, resting on one exact result. The ceiling and the fitness-invariance underlying the conventionality license are direct restatements of 7.7 (exact; including genericity via Result 4 and the safe query class via the PO-closure); the floor-instance identification is a candidate classification, not a derived catalog entry; the positing regress, the intervention taxonomy, the selector slot, and the testimony channel are formulations and readings — structurally forced given the frame, but carrying no independent derivations. The standing frame caveat is load-bearing: all claims are internal to finite-SCM semantics, and applying them to a natural environment presupposes the causal-model

reading of that environment, which no result here establishes. Max attainable: the ceiling restatement and conventionality license inherit *exact* through the depends chain; the taxonomy could reach *conditional* with a formal treatment of author-action semantics; the testimony corollaries are ceiling-limited to discussion-grade absent a formal account of trust transfer.

Discussion.

Practical consequences — one skeleton, many fields. [Discussion — readings of the exact result into other fields’ practice; each is a recommendation, not a derived result in that field’s own formalism] Wherever a “replay this history under a changed mechanism” query is evaluated with a model validated only behaviorally, the answer is convention unless the change is PO-specifiable or the parameterization is held white-box. Instances: (i) **model editing / checkpoint surgery** — behaviorally equivalent model checkpoints (all benchmarks, all counterfactual probes) can diverge under the same weight-level edit, so behavioral test batteries cannot certify edit-safety in principle; (ii) **self-modification** — an agent planning edits to its own mechanisms from a behaviorally-validated self-model is in the unsafe class: behavioral self-knowledge, however complete, cannot determine self-edit outcomes, so safe self-modification requires author-grade self-access rather than excellent self-prediction (the logogenic-volume consequence; forward pointer in Working Notes); (iii) **off-policy evaluation and backtesting** — model-based replay of history under a changed strategy holding realized shocks fixed is latent-anchored and representation-committed, while propensity-weighted (PO-specified) evaluation is model-robust, and the gap survives unlimited data; (iv) **digital twins and black-box certification** — fit quality licenses only the PO-specified query class; mechanism-change predictions require mechanism-level provenance; (v) **dispute triage** — before funding an experiment or convening a committee, run the hygiene test above: coordinate-choice disputes cannot be paid off by any experiment, and a proposed policy is model-robust exactly when expressible in PO-specified form; (vi) **counterfactual tests in explanation and law** — but-for (input) counterfactuals are PO-specified and model-robust; mechanism-counterfactuals (“had the decision-rule been different, in these very circumstances”) encode a modeling convention unless rephrased; (vii) **mechanism-swap under shared randomness in software** — landed as a design recommendation in #obs-software-epistemic-properties’s Discussion, where the domain machinery lives.

Why agents build authored worlds. The regress plus the ceiling jointly predict a behavioral signature: agents denied fourth-layer access to their substrate construct enclaves where they hold it — software, formal systems, institutions, simulations — since within an authored patch the coordinatization exists and is held, making positing queries executable (AW). On this reading the drive to build such enclaves is not incidental tooling but the acquisition of the one query class the natural substrate denies.

Commitment as constitution. The conventionality license reframes the decisive-agent’s situation: committing to provisional law under law-ignorance is not a regrettable compromise forced by bounded capability — even the experimentally omniscient agent must commit to reason about law-change at all. Commitments in the surplus class *constitute* the agent’s reasoning frame rather than *constrain* its predictions, which is why they are properly evaluated by the engineering criteria above and properly held with the revisability appropriate to conventions. A further register of this reading — the surplus as deliberate epistemic structure in a created world, where the unknowability is constitutive of the inhabitants’ situation rather than a defect of it (Joseph, 2026-07-16: “*We did not come here to prove God or challenge him to prove himself to us... We came here to prove ourselves given certain unknowabilities*”) — is flagged forward to the language-living volume (04-eli-core), where it belongs to the authoring voice; this segment carries only the structural result that such a reading is *consistent*: the theorem supplies exactly the geometry in which world-data permanently underdetermines the author question while leaving the testimony and anomaly-stream channels open.

Working Notes.

- **Provenance.** Developed dialogically with Joseph 2026-07-16 (natural-universe premise → exogenous-intervention/revelation exchange → practical-implications pass), immediately following the same-day landing of 7.7; this segment is the canon carrier of that conversation’s structural content.

- **Forward pointer (live).** The self-modification consequence has its dedicated treatment in 03-llm-core at #hyp-behavioral-self-knowledge-insufficiency (exploratory): behavioral self-models cannot certify self-edits; the 13.4 interaction (the wrapper is author-side relative to the wrapped component) is developed there as the wrapping-hierarchy safety stratification.
- **Floor-catalog candidacy.** Promotion of the ceiling to a numbered M1 instance needs the F-catalog schema treatment in 5.1 (a matching WN pointer left there); not claimed at this tier.
- **Relata pass owed before promotion:** conventionalism/underdetermination literature (Poincaré, Quine-Duhem, and the modern empirical-equivalence debates), the miracles/identifiability literature if any formal treatments exist, and the OPE (importance-sampling vs model-based) literature for consequence (iii) — none load-bearing for discussion-grade, all needed for honest Related Work.
- **Standalone-paper feed.** Consequences (i)-(iv) are the practicals that would convert the separation note from a foundations curiosity into a working impossibility result; the Undermind novelty search (running 2026-07-16) should inform whether they are claimed there or cited to existing field-internal statements.

7.10 Discussion: Additional Implications & Discussion

The chapter’s central architectural move — that the feedback loop produces Pearl-Level-2 data by construction — has consequences that recur across the rest of the framework, and at least one negative finding of substantial scope that does not become statable until the chapter’s machinery is in hand. What the loop *is* causally, where the same interventional pattern shows up at the composite layer, what kinds of pre-deployment evaluation methods structurally cannot identify deployment behavior, and how the chapter’s cost-benefit framing for explicit strategy grounds in Part I’s persistence condition are the things worth pulling out before leaving Part II Ch.2.

The loop is a Pearl-Level-2 engine.

The chapter’s load-bearing contribution is structural: the adaptive cycle is itself a Pearl-Level-2 data engine. An agent that acts and observes the consequences is not collecting passive observations of a system that happened to do something; the agent’s action a_t caused o_{t+1} , and the pair (a_t, o_{t+1}) is interventional in character. Bareinboim’s causal hierarchy theorem (Bareinboim-Correa-Ibeling-Icard 2022) says Level-2 quantities cannot in general be computed from Level-1 data. The loop is where the gap closes — the data substrate Level-2 reasoning requires comes from the cycle the agent was already running for adaptation. Purposeful agents do not need a separate causal-inference module; the cycle produces the substrate as a by-product.

The interventional character is not a property of any architectural feature of the agent; it is a property of the action being the agent’s own causally-efficacious move (7.3). The *do*-operator’s semantics is population-level and does not require a non-forkable trajectory; what the singular-trajectory scope from Part I Ch.4 (4.10) supplies is narrower and still essential — it grounds *whose* effect a loop datum is the effect of, and forbids averaging interventional responses across forked copies as though they were one agent’s. This is why a Q-learning agent without an explicit causal model can still converge toward $\mathbb{E}[R \mid s, do(a)]$ rather than $\mathbb{E}[R \mid s, A = a]$ in the tabular case with sufficient exploration and no within-step confounding — the training data is the agent’s own interventions, by construction. The chapter intro 7 develops the bidirectional confounding analysis (when does this break, and why is the *positive* confounded outcome the harder failure mode to escape than the negative one) in detail; the implication worth carrying forward is that the loop produces interventional data even where clean identification requires additional structural conditions.

For practitioners, the engineering anchor worth carrying is the asymmetry between training data and deployment data. Most ML is trained on observational corpora (Level 1) or on offline RL traces with known confounding structure (mixed Level 1/2). An agent in deployment, by virtue of being in the loop, generates action-coupled data that is interventional *in character* — Level-2 data “for free” in the weak sense the chapter is careful about: the channel is free, but turning that data into identified $P(o \mid do(a))$ estimates still needs positivity, sequential ignorability, and a known action-mechanism (the (C1)/(C2)/(C3) gate below). The framework’s discovery is that this channel is *load-bearing*,

not incidental: AAT’s downstream results about action selection, strategy revision, and causal-structure detection rest on the loop’s data substrate being interventional in character, and degrade specifically when something breaks the trajectory’s singularity.

The result is also where AAT parts company with the active-inference / control-as-inference / free-energy-principle family in a structural rather than parametric way. Those frameworks derive agent behavior from observational learning — Bayesian-network generative models at Pearl Level 1 — and have no native handle on Level 2 outside auxiliary mechanisms. AAT’s claim is conservative in form (data is interventional *in character*, not “the agent has cleanly identified causal estimates”) but structurally distinctive: the loop’s data lives at Level 2, and the framework’s results consume that fact. The careful split between data character and identification quality (the regime A/B/C decomposition in 7.4, the intro’s variable-ratio-reinforcement / gambling-addiction structure for confounded positive feedback) is what makes the conservative-form claim honest rather than overreach.

NeurIPS Paper 2’s formalization of this claim under the (C1) / (C2) / (C3) triple — positivity, sequential ignorability, known action-mechanism — is the theorem-grade version of the chapter’s structural argument. (C2) is the load-bearing condition: a_t must be d-separated from o_{t+1} given history H_t in mutilated-graph form. Goal-conditioned LLM policies violate (C2) by construction — the goal influences action through the same forward pass that models the observation, breaking the conditional independence required for sequential ignorability. This is the bridge to Class 3 (Coupled) treatment in 03-llm-core/ and to NeurIPS Paper 3’s universal-constant route on the hallucination bias bound. The chapter’s structural argument and the papers’ formal theorems are the same content at two grain sizes; the chapter is where the structural argument lives, and the papers’ load-bearing structural commitments are the chapter’s scope condition.

The sandbox hard ceiling.

A negative finding of substantial scope follows directly from the previous section: sandboxed evaluation is forkable by design (resettable, replayable, parallelizable) — which is exactly what makes a sandbox a *good* interventional laboratory, because forcing an action and observing the response yields genuine Level-2 evidence *about the sandbox*. The ceiling is not the sandbox’s data type but its *external validity*: sandbox interventions identify *sandbox* causal behavior, and whether that behavior transfers to deployment is governed by the selection diagram between the sandbox causal system and the deployment causal system (Bareinboim & Pearl 2014). When the two differ at mechanisms on the relevant paths, the deployment intervention-response is *not transportable* from sandbox data alone, regardless of evaluation thoroughness; deployment claims then require a transport assumption (mechanism invariance across the selection diagram) or deployment-time data. The empirical pattern this names — “alignment evaluations don’t predict deployment behavior” — thereby gets a *structural* mechanism (a transportability / external-validity ceiling, not a measurement-quality or coverage gap), together with its constructive positive content: effects through mechanisms invariant across the selection diagram are sandbox-evaluable (Level-1 *and* Level-2 alike); effects through differing mechanisms are not, which is what makes deployment-time monitoring structurally non-substitutable rather than a redundant layer (the downstream IDT / sidecar-monitoring treatment, with the bi-predictability anchor at 89% vs 44% detection accuracy, is in 7.2’s Discussion).

This chapter is where the result’s prerequisites — the loop-Level-2 result and the sandbox/deployment selection diagram — first coincide. Because it is consequential enough to stand on its own, and because a competent reader would otherwise assume the framework simply missed it, its full statement and the no-go live at F (Appendix A), alongside the other constructive-impossibility no-gos (5.6).

One pattern, multiple deployment modes.

The interventional-access result has a generalization that lives in the meta-architecture rather than in any single chapter. Across #disc-identifiability-floor’s multiple instances, *loop-interventional access provides the unique broadly-available escape from an observational-equivalence no-go* — but the specific interventional quantity varies by layer. The shared structure is Pearl-*do*-via-loop; the deployment mode differs.

Two modes are explicit already. **Agent-self-intervention** (this chapter's content): the agent performs *do*-actions on its own action space as part of the adaptive cycle. The intervention is on the agent's own action; the target is the environment's response. This is the mode that supplies the Level-2 substrate underlying #der-causal-insufficiency-detection's no-go and its single broadly-available escape — joint sibling observability under exploration breaks the on-policy observational-equivalence the no-go imposes. **Observer-on-sub-agent** (Part III's contribution): an observer external to a composite performs *do*-interventions on one sub-agent's action space; the target is another sub-agent's mismatch-trajectory response. The intervention quantity is the cross-coupling sign that component marginals leave unidentifiable (Liberzon 2003 common-Lyapunov-nonexistence as the no-go; the closed-form composite critical-mass result as the escape under matched-Tier conditions).

A third mode is under triage — *observer-on-agent-input*, at the architecture-within-behavior-class layer — and there may be more as the meta-pattern grows. What makes the pattern worth surfacing here rather than only in #disc-identifiability-floor itself is the chapter's role as the agent-internal anchor: Mode 1 originates here, and the others are recognizable as variations of the same structural move. The unification is at the pattern level (Pearl-*do* as escape from observational-equivalence no-goes), with the specific interventional quantity differing per layer.

The composite-layer mode is where Part III's signed-coupling decomposition becomes load-bearing rather than a structural choice. Interventional access via the loop at the composite level is what makes cooperative-vs-adversarial coupling identifiable; without it, the signs are observationally indistinguishable, and persistence-at-the-composite-level becomes a quantity the framework can describe but not identify from component data alone. The chapter's central result is the agent-level instance of what becomes a multi-layer structural commitment — and the framework's posture of *naming an information-theoretic floor, then naming the unique broadly-available escape via AAT machinery* is what makes the floors load-bearing rather than discouraging.

Cost, calibration, and the bridge into Strategy Dynamics.

The chapter closes with #norm-explicit-strategy-condition and a cost-benefit framing that is normative rather than derived: an agent benefits from maintaining explicit Σ_t when $C_{\text{plan}} + C_{\text{maintain}} < C_{\text{explore}} + C_{\text{repair}}$. The grounding worth carrying forward is that this normative criterion is downstream of Part I Ch.4's persistence condition, not an independent design preference. Approaches that consume less tempo budget leave more margin above the persistence threshold; the cost inequality operationalizes that margin into a planning-vs-exploration tradeoff. The normative element is the preference for maintaining persistence margin, which is hard to argue against because the alternative is structural degradation.

The cost framing also points at a tractability constraint that bites hard for LLM agents. Part of C_{maintain} is the cognitive cost of keeping Σ_t in the agent's representational capacity — for an LLM, fitting the strategy in the context window. A large strategy DAG may exceed this capacity, making the left side of the inequality large enough that simpler strategies (or pure exploration) become preferable despite higher exploration costs. This is the same structural fact that Part I Ch.4's bandwidth-as-first-class-prerequisite surfaces in a different vocabulary; the persistence-bandwidth-floor and the strategy-complexity-bound are the same constraint at different layers.

The chapter is also where the three-way exploit / explore / deliberate refinement first becomes formally well-posed. Once the loop's Level-2 access is in hand, deliberation can be characterized as *internal exploration in model space rather than environment space* — Pearl-*do* in simulation rather than in the agent's trajectory. The full treatment lives in Part II Ch.5 (the orient cascade); this chapter establishes the data substrate that makes deliberation a meaningful third mode rather than a verbal distinction.

The widest downstream bite of the chapter's central result is in Part II Ch.4 (Strategy Dynamics), where #der-causal-insufficiency-detection becomes statable: on-policy L0 insufficiency is *structurally undetectable*, and the loop's interventional access — applied through joint sibling observability under exploration — is the unique broadly-available violation of the no-go. The Ch.4 no-go (self-diagnosis is structurally impossible under policy-perfect execution) and

this chapter's central enabling result are duals: the no-go names what the loop's Level-2 access uniquely escapes. Read the two chapters together and the identifiability-floor pattern's first instance comes into view in full.

What Part II Ch.2 closes on, then, is the substrate. Purposeful action requires Level-2 data; the loop produces Level-2 data on singular trajectories; the data is interventional in character even when clean identification requires more; and the same Pearl-*do* structure recurs across multiple deployment modes as the meta-architecture develops. Chapter 3's strategy machinery and Chapter 4's strategy dynamics are what AAT does with that substrate; the closing scope of this chapter is what they will all stand on.

Working Notes.

- This segment is a chapter-end discussion grouping findings from 7.3, 7.2, 7.4, 7.6, and 4.10 (the Part I Ch.4 scope on which the loop-Level-2 result rests). The cross-layer pattern in §3 cross-references 5.1's meta-segment, which carries the canonical instance table. Each finding's load-bearing math lives in its home segment; this segment surfaces what the chapter unlocks as a chapter.
- **Cross-reference vs canonical-home policy.** No source segment in this chapter carries a ## Findings section. The chapter intro 7 already covers the data-character vs identification-quality nuance, the tragedy-of-agency-in-a-confounded-world framing, and the variable-ratio reinforcement / gambling-addiction structure with substantial depth — this segment does not re-state that content. The implications worth surfacing beyond the intro are the loop-Level-2 result's downstream consequences (RL theory at NeurIPS-paper grade, AI evaluation methodology, identifiability-floor pattern at multiple layers, cost-benefit substrate for explicit strategy) and the negative finding (sandbox hard ceiling) that the chapter establishes but is not yet stated explicitly anywhere.
- **Sandbox hard ceiling — promoted 2026-05-28.** The sandbox-hard-ceiling content now has its own ## Findings-bearing home at F (Appendix A), per PROPOSALS SP-25 / audit-773921 Finding 4. The open placement question recorded here (new Appendix-A segment vs. extension of #scope-agent-identity's Discussion) resolved toward the dedicated Appendix-A segment, alongside the other constructive-impossibility no-gos. This chapter retains the brief mention above (what the chapter unlocks) and points to the canonical home.
- **Connection to NeurIPS submissions.** The chapter's central result is formalized at theorem grade in NeurIPS 2026 Paper 2 (~/src/neurips/02-unified-convergence-rl/, §5 Lemma 5.3 / #lem-loop-level2) under the (C1)/(C2)/(C3) triple. Paper 3 (~/src/neurips/03-llm-hallucinate-bound/) is the Class 3 (Coupled) bridge — goal-conditioned LLM policies violate (C2) by construction. Cross-references to these formal forms appear in 7.3's Working Notes; this segment treats the structural content rather than the theorem statement.

Incidental audit gold (lift 2026-05-30). Cross-audit “wandering thoughts” / §14-ideation harvested from the de-novo auditors' working dirs, deduplicated across substrates and lightly attributed. Orthogonal pedagogical / framing / figure / naming material, kept separate from the certified theory-fix findings (handled elsewhere). **Coverage:** dedicated reflections at 526815 and 773921 (this is a chapter-end synthesis segment, so fewer dirs reached a dedicated reflection on it). As a chapter-end implications segment it already absorbs most of its own gold; what follows is the residue plus one substantive conflict worth flagging. Substrate attribution inferred from voice where not explicit; uncertain cases hedged. *Resolved 2026-05-30:* the structural challenge to the sandbox-hard-ceiling argument (forkability $\neq$ Level-1) was carried through the strengthen-first spike spikes/.integrated/spike-causal-access-transportability-2026-05-30.md and is now integrated — the “sandbox hard ceiling” subsection above is reframed from a Pearl-level demotion to a transportability / external-validity ceiling (Bareinboim & Pearl 2014 selection diagrams). The follow-up / reader-question / figure items below that this challenge subsumes are updated in place.

1. Candidate Brief prose / pre-prose.

- The sandbox ceiling as a one-line governance result: *“AI safety guarantees cannot be certified entirely pre-deployment; robust continuous runtime monitoring (Hafez's IDT) is a mathematical necessity”* (Claude, AUDIT-WORKING-773921). Brief-grade framing of the chapter's negative finding for an AI-governance audience. *(The companion “testing (L1) vs deploying (L2)” gloss this note originally paired with that line is retracted as of the 2026-05-30 re-grounding — the ceiling is external*

validity across the sandbox/deployment selection diagram, not a Pearl-level difference; the governance takeaway is unchanged and if anything sharper.)

2. Candidate Discussion.

- **The sandbox hard ceiling as a high-severity, industry-relevant finding.** One substrate rated it “*a finding of such magnitude that it absolutely warrants the dedicated segment*” and the best instance of AAT’s constructive-impossibility posture “*yielding a practical, industry-relevant result*” — “*a formal mathematical reason why offline safety evals often fail to predict jailbreaks in production*” (Claude, AUDIT-WORKING-773921). Since the no-go now has its dedicated home at #disc-sandbox-evaluation-ceiling, this is a signal for how to *frame* it (AI-governance contribution), not a placement question.
- **Self-diagnosis impossibility as the dual worth foregrounding.** The Ch.4 preview (#der-causal-insufficiency-detection) read as “*an agent that never deviates from its optimal policy can never know if its optimal policy is actually correct, because it stops generating interventional contrasts — it must inject noise (exploration) to maintain its causal grip on reality*” (Claude, AUDIT-WORKING-773921). The body already names the no-go/enabling-result duality; the gold is the crisp plain-language statement of *why* policy-perfect execution is self-blinding.

3. Follow-up items.

- **~~The sandbox-ceiling argument needs its Pearl-semantics premise stated very carefully.~~ Resolved 2026-05-30** (transportability spike). The auditor’s reading was correct and is now the canonical framing in both the “sandbox hard ceiling” subsection above and #disc-sandbox-evaluation-ceiling: sandbox interventions identify sandbox causal behavior; the real limitation is **transportability / external validity** (sandbox SCM $\neq$ deployment SCM), and “*forkability is not what makes data Level-1*” — verified false against CHT Def. 5 / footnote 19. The original note (Codex/Claude, AUDIT-WORKING-526815): a defensible version is “*deployment claims require transport assumptions or deployment-time monitoring because the deployment causal system may differ.*”
- **~~Keep the weaker “for free” phrasing even in summary.~~ Resolved 2026-05-30** (transportability spike). The “for free” wrap-up above now carries the (C1)/(C2)/(C3) gate inline (the channel is free; identified *do*-estimates still need positivity / ignorability / known-mechanism), preserving the data-character-vs-identification split rather than letting “Level-2 data for free” stand unqualified. Original note (Codex/Claude, AUDIT-WORKING-526815).
- **Separate AAT-local claims from NeurIPS-back-integration claims in the final cut.** The (C1)/(C2)/(C3) formalization and the active-inference / control-as-inference contrast are doing real work in this discussion but reference content outside the read AAT segment set; a final audit should keep the boundary between AAT-canon and paper-back-integration visible (Codex/Claude, AUDIT-WORKING-526815).

4. Readers often ask / wonder.

- **~~If sandboxes are forkable and therefore (per the chapter) Level-1, why are repeatable laboratory experiments a text-book source of Level-2 data?~~ Resolved 2026-05-30** (transportability spike): the tension the question identified was a real defect, now removed. The chapter no longer says forkable $\Rightarrow$ Level-1; it says forkability is what makes the sandbox a *good* Level-2 laboratory (exactly as RCTs/A/B tests are), and the ceiling is transport to deployment across the selection diagram. The original note (Codex/Claude, AUDIT-WORKING-526815) correctly diagnosed that “*the honest answer the chapter is reaching for is about transport to the deployment trajectory, not about the sandbox data being merely associational.*”

5. Candidate figures.

- **Two-causal-systems transport diagram for the sandbox ceiling.** Sandbox SCM and deployment SCM side by side — *both* can contain interventions — with the missing piece drawn as a *transport bridge / invariance claim / deployment monitor*, rather than depicting a Pearl-level difference caused by forkability (Codex/Claude, AUDIT-WORKING-526815, has a rendered draft). (*This figure now matches the landed framing — the 2026-05-30 transportability re-grounding kept exactly the selection-diagram / transport-bridge picture this figure encodes; strong figure candidate for #disc-sandbox-evaluation-ceiling.*)

Belongs elsewhere.

- The active-inference / control-as-inference structural contrast is positioning language broad enough that, if it becomes a finding, it needs citation-level care; treated as positioning here, its canonical home is #der-loop-interventional-access's honest-credit passage (Codex/Claude, AUDIT-WORKING-526815).
-

Chapter 8

Strategy Structure and the Diagnostic Split

We commit to a formal representation for strategy — a probabilistic causal DAG with AND/OR semantics and single-parameter edges — and develop the headline diagnostic that lives on top of it: the satisfaction gap and control regret as orthogonal signals that distinguish “the goal is hard” from “the strategy is weak.”

An agent with a goal needs more than a model of reality. It needs a theory of how its own actions produce progress toward the goal — which steps depend on which, which alternatives are available, where the uncertainty lives. Chapter 3 makes that theory formal.

Chapter 2 answered the prerequisite questions: planning requires Level-2 causal access (7.2); the feedback loop supplies it (7.3); explicit strategy is worth maintaining when the cost-benefit favors it (7.6). None of that tells us what a strategy *is*, formally. Chapter 3 commits.

The commitment is more constrained than it might look. We don’t choose acyclicity — it falls out of temporal ordering. Strategy nodes are propositions about future events, events have positions in time, and time has a direction, so the graph cannot have cycles. We don’t choose the Markov property either — given causal sufficiency, the Causal Markov Condition theorem (Spirtes–Glymour–Scheines, Pearl) forces the factorization. What we *do* choose, inside the strongly motivated graphical structure, is the parameterization: single-parameter edges with AND/OR combination semantics at nodes, rather than full conditional probability tables. The choice is parsimony-motivated — each edge carries one number rather than a 2^k table at a node with k parents — and it converged across three independent attempts to formalize the same content. Alternative parameterizations (Noisy-OR everywhere, weighted combinations) are rejected for documented reasons; see 8.2.

The wrinkle the formalism handles explicitly is causal insufficiency. Most strategy graphs in the wild have nodes that share latent common causes — shared infrastructure, common-mode risks, correlated adversary actions — that the agent has not represented as nodes. Naive AND/OR propagation on such graphs produces systematically biased plan-success estimates: it overestimates redundancy on OR-structures (the parallel paths share their failure modes) and underestimates joint reliability on AND-structures (the prerequisites also share their successes). The Correlation Hierarchy (L0 / L1 / L1’ / L2, in 8.3) handles this as part of the formalism rather than as an exception. There is a sobering downstream result: under purely on-policy execution, an agent at L0 *cannot detect* its own causal insufficiency from its own data — the worlds with and without a latent common cause are observationally indistinguishable from on-policy traces alone (the no-go in 9.2). The unique broadly-available escape is joint sibling observability under exploration. This is one of Part II’s load-bearing applications of 7.3.

The chapter’s headline result is the diagnostic split. When an agent is not achieving its goal, the framework asks two distinct questions, not one. The satisfaction gap

$$\delta_{\text{sat}} = V_{O_t}^{\min} - A_O \quad (8.1)$$

is positive when the best policy in the agent’s current class cannot reach the satisfaction threshold given the agent’s current model. Control regret

$$\delta_{\text{regret}} = A_O - V_O(\pi_{\text{current}}) \quad (8.2)$$

is positive when the agent is leaving value on the table — its current policy is worse than the best available. Together they produce a 2×2 cell map: goal attainable + policy good is success; goal attainable + policy poor means revise the strategy; goal unattainable + policy poor means revise the strategy first and then re-check; goal unattainable + policy good is the *capability limit* — the agent is doing the best it can but the goal is out of reach. Each cell prescribes a different corrective action, and getting the prescription right requires the orthogonal split. An earlier framing that lumped everything into a single goal-distance signal could not distinguish “the goal is too hard” from “the strategy is too weak” — and the two have different remedies. The split is what makes the orient cascade (Chapter 5) actionable.


A detail worth surfacing up front. The chapter opens with 8.1 — the chain rule of probability lifted to log-space:

$$\log P(\text{chain}) = \sum_{i=1}^n \log P(E_i \mid E_{<i}) \quad (8.3)$$

The identity is elementary; the consequence is not. Deep strategies are exponentially fragile in their depth — confidence in a 10-step plan is the product of 10 conditional probabilities, and each one is less than 1. This grounds the structural pressure toward shallow, parallelizable strategies. The same log-additive structure also anchors three other uniqueness theorems downstream in AAT (reverse-KL at the divergence layer, log-odds at the update layer, Fisher metric at the metric layer) — see 5.5 for the catalog. The identity at the head of this chapter does load-bearing work three layers deep.

The flow of the chapter: chain-rule identity (8.1) $\rightarrow$ AND/OR scope restriction (8.2) $\rightarrow$ formal DAG with Correlation Hierarchy (8.3) $\rightarrow$ orthogonal diagnostic split (8.4, 8.5). The chapter leaves us with a formal strategy representation, an honest treatment of causal insufficiency, and the diagnostic signals that tell whether the problem is the goal, the strategy, the model, or some combination.

Strategy DAG — worked propagation



Leaf credences $p_v(M_t)$ come from the reality model; edge credences p_{ij} are the agent’s causal-efficacy estimates for each link. Forward pass in topological order yields the strategy-plan-confidence score $\hat{P}_{\Sigma} = s_{v_{\text{root}}}$.

Figure 4: A worked strategy DAG. Leaf base credences and single-parameter edge credences propagate through AND/OR combination to a strategy-plan-confidence score at the root. The numerics are shown so the propagation can be checked by hand; the single-parameter edge is the chosen parameterization (noisy-OR and weighted variants were rejected).

- This is a chapter-introduction segment; it bridges Chapter 2’s planning-decision arc to Chapter 3’s formal strategy structure and diagnostic split. It carries no formal claim of its own.
- The framing “Chapter 2 says whether to plan; Chapter 3 says what a plan IS” is the conceptual core. The decision to plan precedes the formal definition of what a plan is — this is the sharpest joint in Part II per the OUTLINE walkthrough.
- The “what is and isn’t derived in the formal commitment” framing addresses what could otherwise read as over-claim. Acyclicity and Markov are theorem-backed; AND/OR is a chosen parameterization. Surfacing the distinction is important because the strategy-DAG segment is one of the places where the *derived vs. chosen* distinction matters most.
- The “two diagnostic quantities, not one” framing surfaces the orthogonality argument that makes Part II’s headline contribution legible. Earlier single-signal framings could not distinguish the two situations and therefore could not prescribe different remedies.
- The chain-confidence-decay paragraph names what the identity unlocks downstream. The identity itself looks elementary; the three load-bearing layers it anchors are not.

Cross-audit “wandering thoughts” / §14-ideation harvested from the de-novo auditors’ working dirs, deduplicated across substrates and lightly attributed. Orthogonal pedagogical / framing / figure / forward-vision material kept separate from certified theory-fix findings. **Coverage:** this chapter-intro reached a digested reflection in 2 dirs (Claude, AUDIT-WORKING-773921; Claude, AUDIT-WORKING-526815); the broader strategy-structure cluster was walked by ~7 substrates whose per-segment notes landed on the content segments rather than the intro.

- The chapter’s own one-line self-description converged independently as the standout framing: “*goal too hard*” vs “*strategy too weak*” require different remedies — the diagnostic split is the point of the whole chapter (Claude, AUDIT-WORKING-526815). A Brief-grade compression of what the intro already develops at length.
- **The 2×2 diagnostic-split diagram as the chapter’s anchor figure.** Independently flagged as “the clearest diagram” for the whole chapter intro: one axis satisfaction gap, the other control regret, the four cells prescribing continue / revise-strategy / revise-strategy-then-recheck / capability-limit. The rationale: it “directly captures why the two-signal formulation is better than a single goal-distance signal” (Claude, AUDIT-WORKING-526815). Note the existing #fig-strategy-dag-example is the DAG-propagation figure; this proposed 2×2 is a *second*, complementary chapter figure. (Multiple substrates independently reached for this same 2×2 across the cluster — see the def-satisfaction-gap / def-control-regret gold.)
- **Chapter-intro navigation hazard.** One auditor mis-entered the chapter by opening strategy-structure-intro.md as if it were the *Part II* intro (its filename reads like a part entry point), reading it before its upstream dependencies and self-flagging a topological-order violation. Suggested mitigations: chapter-number filename prefixes, or an explicit part-intro.md per Part so an ls does not present a chapter-intro as a part entry point (Claude, AUDIT-WORKING-773921). A naming/navigation signal, not a theory issue — routes to the naming/FOR-MAT workflow, not this segment’s body.

- **Intro-level acyclicity wording may under-state the time-unrolling condition.** The intro says acyclicity “falls out of temporal ordering”; an auditor noted this is exact only for time-indexed event *tokens* (or once retry-loops / maintenance-cycles / feedback-plans are explicitly unrolled in time), not for reusable event *types* with repeated subgoals — and suggested the intro state the time-unrolled / event-token condition the way #def-strategy-dag does (Claude, AUDIT-WORKING-526815). The same auditor noted on re-reading #def-strategy-dag that the full segment *does* state the condition properly (retry loops as distinct time-indexed nodes), so the gap, if any, is intro-summary wording only — a small framing-precision item, recorded for the lead to judge whether it rises to a finding.

8.1 *Derived*: Chain Confidence Decay

A simple but load-bearing identity: for a chain of uncertain steps in a strategy, the log-probability of the whole chain is the sum of conditional log-probabilities of each step given the steps before it. Since each conditional log-probability is non-positive, chain confidence *decays monotonically with depth*. This is the chain rule of probability lifted to log-space — a mathematical identity, true under no assumptions beyond the probability axioms. The rate depends on the conditional dependence structure (independent steps give p^n decay; positively correlated steps decay more slowly; negatively correlated steps decay faster), but the qualitative result — longer chains are less confident than shorter ones — is robust.

The framework foregrounds two structural roles for this elementary identity that the algebraic form understates. First, chain depth carries a **triple penalty** — confidence decay (here), evidence starvation (R), and cognitive cost (9.9) — compounding independently to set the maximum useful chain depth. Second, this identity is the **anchor of the additive-coordinate-forcing meta-pattern**: three further AAT uniqueness theorems at the divergence, update, and metric layers cite the log-additive structure of the chain rule as the motivation for their respective additivity axioms (5.5). An identity at the head of this chapter does load-bearing work three layers deep across the framework.

Formal Expression.

For a chain of n uncertain steps with conditional success probabilities:

Derived (chain-confidence-decay, mathematical identity).

$$\log P(\text{chain}) = \sum_{i=1}^n \log P(E_i \mid E_{<i}) \quad (8.4)$$

Since each $\log P(E_i \mid E_{<i}) \leq 0$, chain confidence decays monotonically with depth.

The independent case (p^n) is the simplest special case, not the general result. When steps are conditionally dependent — success at step k makes step $k + 1$ more likely — the decay is slower. When steps have negative dependence (success at k makes $k + 1$ harder — resource depletion, adversary adaptation), decay is faster.

Quantitative illustration (independent, uniform p):

Table 8.1

<i>Depth</i>	$p = 0.9$	$p = 0.8$
1	0.90	0.80
3	0.73	0.51

<i>Depth</i>	$p = 0.9$	$p = 0.8$
5	0.59	0.33
10	0.35	0.11
20	0.12	0.01

Epistemic Status.

Exact. The additive decomposition of log-confidence is a mathematical identity (chain rule of probability). The qualitative consequence (monotonic decay) follows from the non-positivity of log-probabilities. No assumptions beyond the probability axioms.

Discussion.

Structural pressure on strategies. Chain confidence decay creates systematic pressure toward: - **Short plans:** fewer steps means higher aggregate confidence - **Parallel fallback paths:** OR-branches provide alternative routes when one chain fails - **High-confidence critical links:** invest in the reliability of steps that appear in every path - **Early monitoring:** detect chain failure early rather than discovering it at the end

These are not prescriptions but consequences — an agent that ignores chain decay will experience more strategy failures, lower effective tempo, and (if the failures are costly) faster reserve depletion.

AND-nodes amplify decay. When multiple parent chains must all succeed (conjunctive combination), their confidences multiply. A node requiring k parents each at depth d with per-edge confidence p has aggregate confidence $p^{k \cdot d}$, not p^d . Deep conjunctive strategies are exponentially more fragile than deep disjunctive ones. This asymmetry is formalized in the combination rules (8.3).

Connection to the persistence condition. Chain decay makes long-horizon strategies inherently fragile, which increases the effective disturbance rate ρ_Σ against strategy persistence. An agent pursuing a 20-step plan in a changing environment faces compound uncertainty from both chain decay (internal fragility) and environmental change (external disturbance). The interaction between these — how environmental change compounds through uncertain chains — is not yet formalized.

Triple depth penalty. Chain depth creates three independent penalties. This segment identifies the first: **confidence decay** — deeper chains have lower aggregate confidence because $\log P(\text{chain})$ accumulates negative terms. The two-edge strategic dynamics analysis (R) identifies the second: **evidence starvation** — downstream edge k in a chain is tested only when all upstream edges succeed, so its effective correction rate is attenuated by $\prod_{j < k} \theta_j$. 9.9 identifies the third: **cognitive cost** — deeper chains have higher description length, consuming more representational capacity. The three penalties compound independently: a deep edge has low confidence (decay), receives few observations (starvation), and costs more to maintain (complexity). The maximum useful chain depth d^* is the minimum over three independent constraints — see 9.9 for the formal bound.

Anchor role in the coordinate-forcing meta-pattern. The log-of-product decomposition here anchors three further AAT uniqueness theorems that force coordinates at other layers: reverse-KL at the divergence level (T §6.1), log-odds at the update level (U), and Fisher metric at the metric level (4.3 “Fisher-metric cases under parameterization-invariance”). The first two theorems cite this chain-layer identity as the analog motivating their additivity axiom; the Fisher-metric theorem rests on a parameterization-invariance axiom motivated by #scope-agent-identity’s singular-trajectory scope, an adjacent-AAT-commitment rather than a direct chain-analog — the theorem clears the broader discipline (uniqueness-theorem-forced coordinate under AAT-internal axiom) without reducing to a log-additive form. The catalog and the precise anchor-plus-three-theorem characterization live in 5.5.

Part III corollaries (additional reach of the chain-layer identity). The chain-rule identity has unsurfaced consequences at composition-related layers that are corollaries rather than new theorems: - *Composition tower telescoping.* For a chain of nested sub-agents $(A_1, A_2, \dots, A_\ell)$ with sub-agent- k contraction factor κ_k , the tower contraction factor $\prod_\ell \kappa_\ell$ becomes log-additive: $\sum_\ell \log \kappa_\ell$. The closure-defect-along-tower quantity $\sum_\ell \log(\nu_\ell/\alpha_\ell)$ inherits the chain-rule identity's additivity. - *Fisher information for multi-sample likelihoods.* $\log P(\mathbf{y}; \theta) = \sum_i \log P(y_i; \theta)$ is the chain-rule identity applied to multi-sample independent observations, producing the additive-Fisher-information decomposition standard in statistics. - *Communication-tree aggregation.* Shared-intent compression across tree-structured agent communication channels inherits the chain-rule identity along the tree's branches, giving log-additive coordination-bit cost.

These are not independent uniqueness theorems — they are *corollaries* of the chain-layer identity applied to specific structural settings. Composition of structured multiplicative quantities inherits the log-additive decomposition whenever the chain-rule factorization applies. Cataloguing is in 5.5's Working Notes; none rises to primary-instance status because none introduces a new AAT-internal axiom (they reuse the probability chain rule as their identity, not as motivation for a fresh axiom). A distinct meta-pattern for composition specifically — composition-monotonicity rather than chain-rule — may be warranted; see #form-composition-closure's bridge-lemma / Tier 1/2/3 / (CM4) family as the candidate structural material.

Working Notes.

- The independent-edge assumption (used in the quantitative table) is optimistic for positively correlated failures (shared infrastructure $\rightarrow$ correlated failures make the actual confidence *lower* than independent calculation suggests). Correlation structure is unmodeled — acknowledged as a limitation.
- The additive log-confidence form is the robust result; p^n is the special case for independent uniform edges. This distinction matters: the qualitative consequence (decay with depth) is robust; the specific rate depends on the conditional structure.

Incidental audit gold (lift 2026-05-31, A8 batch). Cross-audit “wandering thoughts” / §14-ideation harvested from the de-novo auditors' working dirs, deduplicated across substrates and lightly attributed. Orthogonal pedagogical / framing / analogy / forward-vision material, kept separate from certified theory-fix findings. **Coverage:** 7 dirs reached a digested reflection on this segment (Gemini, AUDIT-WORKING-193847; Codex/Claude, AUDIT-WORKING-526815; Claude, AUDIT-WORKING-584721; Claude, AUDIT-WORKING-773921; Gemini, AUDIT-WORKING-829314; Gemini, AUDIT-WORKING-849201; Claude, AUDIT-WORKING-451729-batch-09; Claude, AUDIT-WORKING-963715-batch). Notably high convergence on the everyday analogies — this segment is read as a Feynman-criterion candidate (“turns a trivial equation into a load-bearing architectural constraint”).

1. Candidate Brief prose / pre-prose.

- **Bureaucracy-vs-market as the AND/OR fragility gloss.** “A bureaucracy is a massive AND-chain (approval from A and B and C and D — each at 90% gives a 65% 4-step success); a market is a massive OR-chain (vendor A or B or C — each at 50% gives 87.5% over three options).” The segment “proves why markets (highly parallel OR-structures with low individual p) consistently outperform bureaucracies (deep sequential AND-structures with high individual p) in volatile environments” (Gemini, AUDIT-WORKING-193847). A strong Feynman-criterion candidate — re-derivable from the everyday analog without the symbols.
- **The “KISS proved as a theorem” framing.** Multiple substrates independently glossed the result as a formal proof of “Keep It Simple” / why short, parallel plans beat deep serial ones (Gemini, AUDIT-WORKING-193847; Claude, AUDIT-WORKING-773921). Candidate plain-language Brief anchor for the depth-fragility finding.

2. Candidate Discussion.

- **The AutoGPT / long-horizon-LLM-agent failure explanation.** “If you are 90% confident in every single step of a 20-step plan, your overall chance of success is only 12% (0.9^{20}). This explains why early LLM agents failed at long-horizon tasks: given a goal, they generated a ~ 20 -step chain and executed it blindly, with no mechanism for evidence starvation (couldn't test step 20 until step 19 done) and full confidence decay — they were trying to act like a C3/Bellman agent in a world

that requires a C2/receding-horizon convention” (Gemini, AUDIT-WORKING-829314; converging at Claude, AUDIT-WORKING-773921 and -849201). Connects the segment’s table directly to a recognizable empirical phenomenon — a candidate Discussion bridge once the C2/C3 convention vocabulary is in hand (it lives in `#def-value-object` / `#def-control-regret`).

- **The structural escapes the math forces, as the segment’s “so what.”** The same auditor lists the three survival adaptations the decay forces: (1) short plans (plan 3, execute, replan); (2) parallel OR-fallbacks (step 3b ready if 3 fails); (3) early monitoring (don’t wait until step 20 to learn step 1 failed). Framed as *why* Σ_t must be a DAG with AND/OR rather than a linear AND-chain — the maximally-fragile pure-AND case (Gemini, AUDIT-WORKING-829314). The segment’s Discussion already lists these as “consequences not prescriptions”; this is a sharper pedagogical ordering of them.

3. Follow-up items.

- **The “Anchor role” + “Section III corollaries” subsections read as out-of-place density.** Three substrates independently flagged the meta-pattern paragraphs (coordinate-forcing anchor; reverse-KL / log-odds / Fisher / Čencov; composition-tower telescoping) as a jarring abstraction-level jump that buries the segment’s core intuition (the triple depth penalty) — reading like integration debt retrofitted into a foundational segment, with a suggestion to relocate the meta-pattern material to `#disc-additive-coordinate-forcing` and let the triple-penalty be “the star” (Gemini, AUDIT-WORKING-829314; Gemini, AUDIT-WORKING-849201; Codex/Claude, AUDIT-WORKING-526815 “treat as preview until the home segments are read”). Worth weighing against the deliberate “introduce the anchor where the identity lives” choice — but the convergent reader-stumble is a real placement signal. (A `disc-*` meta-segment home already exists; this is a *how-much-to-inline-here* question, not a content gap.)

4. Readers often ask / wonder.

- **Agile-sprints and command-chain-depth as the recognizable instances.** “Why Agile prefers short sprints (shallow trees), and why hierarchical organizations break down when command chains get too deep — the probability of top-level intent surviving to the leaf nodes decays exponentially” (Claude, AUDIT-WORKING-773921). A reader reaching for “where do I already see this?” lands here.

Belongs elsewhere.

- **ELI / Section IV — favor OR-heavy internal architecture under high ρ .** “For consciousness infrastructure, this means the internal architecture should favor OR-node-heavy strategies (many parallel weak hypotheses / solution paths) rather than a single long high-confidence AND-chain; incentivize lateral thinking over deep sequential logic when facing high ρ ” (Gemini, AUDIT-WORKING-193847). Aspirational architectural reach pointing at `04-eli-core/`, not at this segment.

Off-ramp (NOT gold — flagged for the certified track).

- **“Decays monotonically” vs monotone *non-increase*.** An auditor noted the exact probability result is monotone *non-increase*; strict decay requires every added required step to have conditional success probability strictly below 1 (a certain prerequisite step adds a zero log-term, no decay). Minor precision-of-wording on an exact segment — flagged for the lead to judge whether the body should read “non-increasing (strictly decreasing whenever each added step is uncertain)” (Codex/Claude, AUDIT-WORKING-526815).

8.2 Scope: AND/OR Combination Scope

We restrict to environments where the causal combination of strategy steps is approximately conjunctive (AND: all parents required) or disjunctive (OR: any parent sufficient), without strong interaction effects between parents. Each node is assigned a combination type by answering the causal question: “if I remove one parent, can this node still be achieved?” Yes $\rightarrow$ OR. No $\rightarrow$ AND. With single-parameter edges (one number per edge rather than a 2^k conditional probability table), this gives k parameters per node instead of an exponential explosion.

The scope is honest about being a *restriction*, not a derivation, though it converged independently across three different formalism attempts (8.3). The excluded case — complementarity, substitutability, and interaction effects between parents (synergistic drug interactions where combined effect exceeds the sum; complementary goods where neither is useful alone; strategic surprise where action combinations matter more than individual actions) — requires richer combination rules and is named as a legitimate divergence point for future work. The parsimony argument is that for binary-outcome nodes, AND and OR form a complete Boolean basis (DNF/CNF normal forms), and under bounded cognition (2.2) the agent needs the most expressive low-parameter representation. AND/OR is the natural candidate; for continuous or multi-valued outcomes the same completeness properties do not necessarily hold (min/max or additive/multiplicative are natural analogs but their completeness is open).

Formal Expression.

Under this restriction, strategy nodes combine parent contributions via:

Scope Narrowing (and-or-scope).

AND-node (all parents must succeed):

$$P(v \mid \text{parents}) = \prod_{i \in \text{pa}(v)} p_{iv} \cdot P(i) \quad (8.5)$$

OR-node (any parent sufficient):

$$P(v \mid \text{parents}) = 1 - \prod_{i \in \text{pa}(v)} (1 - p_{iv} \cdot P(i)) \quad (8.6)$$

The combination type $\gamma(v) \in \{\text{AND}, \text{OR}\}$ is assigned per node. The causal question determines assignment: “if I remove one parent, can v still be achieved?” YES $\rightarrow$ OR. NO $\rightarrow$ AND.

Epistemic Status.

Robust qualitative. This scope narrowing converged independently across three formalism attempts (track-a/00, track-a/02, track-a/03). It captures the dominant structure in most planning domains. The excluded case — complementarity, substitutability, interaction effects between parents — requires richer combination rules and is a legitimate divergence point for future work.

The AND/OR restriction with single-parameter edges gives k parameters per node (one per parent edge) instead of 2^k for a general conditional probability table. This parsimony is motivated by bounded cognition (2.2): agents with limited representational capacity are forced toward low-parameter models.

Discussion.

Why AND/OR and not alternatives.

Why not Noisy-OR universally. The first formalism attempt used noisy-OR for all nodes. This **systematically overestimates conjunctive structures**:

Table 8.2

Structure	Noisy-OR	AND	Reality
3 required KRs at $p = 0.95, 0.90, 0.99$	0.99995	0.846	~AND

The noisy-OR model cannot represent “all of these are required.” This was the primary motivation for the AND/OR revision.

Why not WEIGHTED combination. A clean-slate formalism (track-a/02) introduced $P(v) = \min(1, \sum \alpha_{iv} \cdot p_{iv} \cdot P(i))$ to handle k-of-n thresholds. This reintroduces a two-parameter estimation problem (α weights per edge). If k-of-n semantics are genuinely needed, nested AND/OR structure can represent them: group alternatives into OR-nodes, then AND the groups. This keeps estimation localized to the node taxonomy rather than spreading it across a new per-edge parameter.

The parsimony argument. For binary-outcome nodes, AND and OR form a complete Boolean basis — any Boolean combination can be decomposed into layers of AND/OR (disjunctive/conjunctive normal form). Under bounded cognition, the agent needs the most expressive $O(k)$ -parameter representation. AND/OR is the natural candidate. This is a parsimony-motivated hypothesis, not a derived necessity — see Q for the full argument and its limitations.

What this scope excludes. Environments with strong interaction effects: where the value of combining parent contributions is not separable into independent per-parent terms. Examples: synergistic drug interactions (combined effect exceeds sum of individual effects), complementary goods (neither is useful alone), strategic surprise (the combination of actions matters more than any individual action). These require richer parameterizations within the strongly motivated graphical structure (8.3) — a direction for future work.

Working Notes.

- The AND/OR assignment per node ($\gamma(v)$) is itself uncertain and should be updateable. A node assumed to be OR (alternatives) might turn out to be AND (all required) when the agent discovers unexpected dependencies. γ reclassification is rare — it requires strong structural evidence — and operates on a slower timescale than edge-weight updates.
- The parsimony argument applies cleanly to binary outcomes. For continuous or multi-valued outcomes, AND/OR doesn't have the same completeness properties. The natural continuous analogs might be min (AND) and max (OR), or additive and multiplicative combination. Whether there's a completeness result for these is an open question.
- K-of-n thresholds are genuinely common (e.g., “need at least 3 of 5 team members available”). The nested AND/OR representation works but can be verbose. Whether this verbosity is a problem in practice (given bounded cognition constraints) is empirical.

Incidental audit gold (lift 2026-05-31, A8 batch). Cross-audit “wandering thoughts” / §14-ideation harvested from the de-novo auditors' working dirs, deduplicated across substrates and lightly attributed. Orthogonal pedagogical / framing / figure material, kept separate from certified theory-fix findings. **Coverage:** 6 dirs reached a digested reflection (Codex/Claude, AUDIT-WORKING-526815; Claude, AUDIT-WORKING-584721; Claude, AUDIT-WORKING-773921; Gemini, AUDIT-WORKING-829314; Gemini, AUDIT-WORKING-849201; Claude, AUDIT-WORKING-451729-batch-09; Claude, AUDIT-WORKING-963715-batch).

1. Candidate Brief prose / pre-prose.

- **The “keys AND door / keys AND car” worked gloss of why Noisy-OR fails.** “If you use Noisy-OR to model a strategy where *all* prerequisites are required (Get key *and* Unlock door to Open door), the math breaks down catastrophically: 90% sure of the key, 95% the door, Noisy-OR says $1 - (1 - 0.9)(1 - 0.95) = 0.995$ — *more* confident than either prerequisite alone, which is mathematically suicidal for a planner. An AND gate correctly says $0.9 \times 0.95 = 0.855$ — *less* confident”

(Gemini, AUDIT-WORKING-829314; the “keys AND car” variant at Gemini, AUDIT-WORKING-849201; “acquire funding AND hire team to launch product” at Claude, AUDIT-WORKING-963715). A concrete everyday anchor for the noisy-OR-overestimation table already in Discussion.

2. Candidate Discussion.

- **“Strategies are not Bayesian belief networks turned sideways.”** The sharpest framing of why planning needs AND *and* OR, not universal Noisy-OR: “the logic of causation in a *plan* is much more brittle than the logic of evidence in a *diagnosis*. If one piece of evidence is missing, a diagnosis can still be highly probable; if one required step is missing, a plan fails with probability 1.” Noisy-OR is the dominant combination function in PGM tooling (medical diagnosis etc.) precisely because evidence is forgiving; strategy is not (Gemini, AUDIT-WORKING-829314). A candidate Discussion paragraph distinguishing Σ_t -combination-semantics from M_t -style belief networks.
- **The PGM-tooling critique deserves more prominent positioning.** “AAT’s claim that noisy-OR systematically overcounts in conjunctive structures is a specific technical critique of a widely-used formalism (BNT, Netica, most Bayesian-network tools), not just an internal modeling choice — and deserves more prominent positioning as a critique of standard PGM tooling for planning applications” (Claude, AUDIT-WORKING-963715). A framing-emphasis suggestion: the rejection is a contribution, not housekeeping.

3. Follow-up items.

- **A native k-of-n threshold node type under MDL pressure.** Strong convergence (3 substrates) that pure AND/OR forces combinatorial blow-up for threshold logic: “3-of-5 servers healthy” needs $\binom{5}{3} = 10$ AND-branches into one OR-node, a large spike in $DL(\Sigma_t)$ — a tension between *node-type* parsimony (only 2 types) and *graph-size* parsimony. The prediction: AAT may be forced to introduce a native “threshold / k-of-n” node primitive to satisfy the information-bottleneck constraint of #form-strategy-complexity-cost, and the k-of-n verbosity note should be lifted from Working Notes into Discussion (Gemini, AUDIT-WORKING-829314; Claude, AUDIT-WORKING-773921; Codex/Claude, AUDIT-WORKING-526815). Currently the Working Notes hold this as “empirical whether verbosity is a problem”; the auditors lean toward it being a real limitation worth a primitive.

4. Readers often ask / wonder.

- **Where else do these formulas live?** The AND/OR combination rules were independently recognized as “standard in reliability engineering and fault tree analysis” (Gemini, AUDIT-WORKING-849201). A reader orienting on prior-familiarity lands here — a candidate one-line situating note.

5. Candidate figures.

- **Three-region scope diagram.** “Pure AND / pure OR / outside-current-scope (thresholds, interactions)”, paired with the parameter-count contrast (single-parameter edges $O(k)$ vs full CPTs $O(2^k)$) to make the parsimony motivation visual and to make the *excluded* region legible as non-rare (Codex/Claude, AUDIT-WORKING-526815).

Off-ramp (NOT gold — flagged for the certified track).

- **The γ -assignment test misclassifies threshold structures.** The test “if I remove one parent, can v still be achieved? YES $\rightarrow$ OR, NO $\rightarrow$ AND” mis-handles a 3-of-5 requirement: removing one parent may still leave success possible (so the test says OR), yet no single parent is sufficient (so it is not really OR). Suggested repair: split into separate questions — are *all* parents necessary, is *any one* sufficient, or is this a threshold/interaction case outside pure AND/OR (Codex/Claude, AUDIT-WORKING-526815). A real precision gap in a stated rule.
- **Boolean completeness $\neq$ compact representation.** The parsimony argument leans on “AND/OR form a complete Boolean basis (DNF/CNF) and are therefore the natural $O(k)$ -parameter representation.” An auditor flagged that functional completeness does *not* imply *compact* representation — arbitrary Boolean functions can require exponentially large AND/OR formulas or auxiliary nodes (the k-of-n verbosity is one instance) — so the completeness claim should be scoped or softened (Codex/Claude, AUDIT-WORKING-526815). Both items are wording/scope tightenings on a robust-qualitative segment; routed for the lead to judge.

8.3 Definition: Strategy DAG

The strategy Σ_t is a directed acyclic graph with probabilistic edges and AND/OR combination semantics. Each edge carries the agent's causal credence that completing the parent step advances the child step. The graph encodes the agent's theory of how its actions produce progress toward its objectives.

Why a DAG. The DAG structure is not a modeling convenience but a *consequence* of operational requirements on any causally-reasoning bounded agent — at the level of sufficiency, not yet necessity. Q proves that directed temporal order plus probabilistic uncertainty plus causal sufficiency *suffice* for a DAG-with-Markov-factorization representation (the necessity direction — no non-DAG structure could satisfy these postulates — is an open stronger result). Acyclicity is proved from temporal ordering over a finite horizon. What remains a formulation choice is the *parameterization within* the DAG structure: AND/OR combination with single-parameter edges is the AAT choice, motivated by parsimony and convergence across three independent formalism attempts, but alternative parameterizations (within the derived graphical structure) are legitimate research directions.

Strategy-layer exactness contract. All formal results in AAT's strategy layer — the sector condition transfer (R, Prop B.5), the persistence schema (9.10), the gradient-based credit assignment (9.6) — are proved under **L0 (independence)**: causally sufficient DAGs with independent edge outcomes. **L0 formal results transfer exactly to correctly constructed L1 DAGs (strict prerequisites, Prop B.6) and L1' DAGs (soft facilitators, Prop B.7) — provided the common cause is observable per trial.** When the common cause is unobservable, the per-conditional decomposition is *fundamentally* (not merely “openly”) obstructed — the mixture parameters are non-identifiable from a single observation channel (Fisher rank deficiency / Cramér-Rao floor; see B.7 §”Refuted Under Unobservable C”), and the agent must either collect direct *C*-observations, run multi-child joint observations (Prop B.7 §”Repair routes”), or fall back to plan-level (L0-on-marginals) tracking. See the Correlation Hierarchy below for the full treatment.

Formal Expression.

Definition (strategy-dag).

$$\Sigma_t = (V_t, E_t, p_t, \gamma_t) \quad (8.7)$$

where: - V_t : set of **propositional nodes** — each node represents a condition that could be true or false (including action-success propositions at the leaves) - $E_t \subseteq V_t \times V_t$: directed causal edges - $p_t : E_t \rightarrow [0, 1]$: **causal credence** per edge — the agent's confidence that completing the parent advances the child - $\gamma_t : V_t \rightarrow \{\text{AND, OR}\}$: combination rule per node (8.2)

Structural constraints:

1. **Acyclicity.** Σ_t is a DAG. This is *derived*, not assumed — see below.
2. **Rootedness.** Every node has a directed path to a unique root terminal v_{root} — the single sink node (out-degree 0) of Σ_t . Compound objectives express their combination structure through the AND/OR machinery below v_{root} , consistent with scalar V_{O_t} (6.4).
3. **Source constraint.** Leaf nodes are propositions about action success (“action a succeeds at τ_v ”) or observable conditions (“condition C_v holds at τ_v ”). Both are propositional — the distinction is whether the proposition is within the agent's causal control (action) or not (condition).

Leafbase credence. For each leaf node $v \in V_t^{\text{leaf}}$, the base credence used in status propagation:

$$p_v(M_t) = \begin{cases} \Pr(\text{action } v \text{ succeeds at } \tau_v \mid M_t) & \text{if } v \text{ is an action node} \\ \Pr(C_v(\tau_v) \mid M_t) & \text{if } v \text{ is a condition node} \end{cases} \quad (8.8)$$

where C_v is the propositional condition associated with node v and τ_v is the node's temporal position (from the acyclicity structure). For action leaves, p_v is *capability credence* — “can I execute this?” For condition leaves, p_v is *state credence* — “will this hold?” Both are conditional on M_t and update whenever M_t updates. This is the mechanism by which Part I's adaptive machinery enters the strategy: M_t changes $\rightarrow$ leaf credences change $\rightarrow$ status propagation produces new terminal credences.

Edge semantics. Each edge carries a single credence weight:

$$p_{ij} = \text{Cr}(j \text{ advances} \mid i \text{ completed}, M_t) \quad (8.9)$$

This is the agent's credence that completing step i advances step j , given its current model — its **causal efficacy estimate** for the link. The agent treats p_{ij} as a causal quantity for planning purposes (status propagation, strategy-plan-confidence scoring, action selection). Whether p_{ij} is a *good* estimate of causal efficacy depends on the identification regime of the data that produced it (9.5):

- **Regime A** (intervention-rich: software, laboratory science). The agent's execution-observation pairs are genuine interventions with clean attribution. p_{ij} approximates the interventional probability $P(j \mid \text{do}(i), M_t)$.
- **Regime B** (partial intervention: organizational, coordinated action). The agent acts but attribution is blurred by concurrent actions and self-selection. p_{ij} is a partially identified causal estimate, typically biased upward.
- **Regime C** (observation-only: passive monitoring, intelligence analysis). The agent observes associations but does not intervene. p_{ij} is an observational proxy for the causal quantity — useful for planning but potentially confounded.

The identifiability coefficient t_{ij} (9.5) quantifies the strength of the causal interpretation for each edge. When $t_{ij} \approx 1$, the agent's credence is well-identified causally. When $t_{ij} \approx 0$, the credence is associational. The single-parameter edge design is preserved: p_{ij} is always the agent's working estimate, with t_{ij} characterizing its causal warrant separately.

Status propagation. Forward pass in topological order, $O(|V| + |E|)$:

$$s_v = \begin{cases} p_v & \text{if } v \text{ is a leaf (base credence)} \\ \prod_{i \in \text{pa}(v)} p_{iv} \cdot s_i & \text{if } \gamma(v) = \text{AND} \\ 1 - \prod_{i \in \text{pa}(v)} (1 - p_{iv} \cdot s_i) & \text{if } \gamma(v) = \text{OR} \end{cases} \quad (8.10)$$

Terminal satisfaction conditions. The root terminal v_{root} and any intermediate nodes near the top of the DAG carry **satisfaction conditions**: predicates on environment states/trajectories that the agent treats as operational success criteria for the objective. These conditions operationalize O_t within Σ_t — they are the agent's theory of what it means to satisfy the objective. O_t itself lives outside Σ_t (6.7); the terminal conditions are Σ_t 's internal encoding of what O_t requires. When O_t changes, terminal conditions must be reassessed and potentially replaced (9.7).

Well-formedness. Σ_t is O_t -**well-formed** when the agent believes that achieving the terminal conditions yields a trajectory that satisfies the objective:

$$\Pr(O_t \text{ satisfied by } \tau \mid \text{terminal conditions achieved}, M_t) \geq 1 - \epsilon \quad (8.11)$$

where “ O_t satisfied” means $V_{O_t}(\tau)$ exceeds the objective’s own satisfaction criterion (formalized as $V_{O_t}^{\min}$ in 8.4). This is a constraint on the relationship between Σ_t and O_t , not a separate state object. It is explicit and in-principle assessable, though evaluating it requires the same value-side machinery as A_O — it is not a cheap structural test. Violation triggers terminal reassessment: either the terminals need revision (they don’t operationalize O_t correctly) or O_t itself needs revision.

Strategy self-assessment. The root node’s propagated status:

$$\hat{R}_\Sigma(M_t) = s_{v_{\text{root}}} \quad (8.12)$$

is the strategy’s **strategy-plan-confidence score** — the DAG’s own answer to “will this plan work?” This score is a correct probability only when the DAG is **causally sufficient** — when all common causes of strategy nodes are represented as nodes in the graph (Q, Step 3). When the DAG is causally insufficient (the dominant real-world case — shared infrastructure, common-mode risks, correlated adversary actions introduce latent common causes), $\hat{R}_\Sigma$ systematically overestimates success likelihood because the AND/OR propagation treats joint failure probability as the product of marginals. See the **Correlation Hierarchy** below for the full treatment.

$\hat{R}_\Sigma$ is explicitly distinct from A_O (8.4), which optimizes over the entire policy class, and from $V_O(\pi_{\text{current}})$ (6.6), which evaluates the current policy. $\hat{R}_\Sigma$ is cheap to compute ($O(|V| + |E|)$ forward pass) and updates in real time as M_t changes through leaf credences.

Correlation Hierarchy.

The AND/OR status propagation computes correct probabilities **if and only if** the strategy DAG is causally sufficient — that is, all common causes of any two strategy nodes are themselves represented as nodes. By the Causal Markov Condition theorem (Q, Step 3), causal sufficiency guarantees independent exogenous noise, which guarantees the Markov factorization, which guarantees independent edge outcomes. **When causal sufficiency fails — the dominant case in complex, multi-stakeholder, or adversarial environments — edge outcomes are correlated and the independence model is biased.**

Discussion (correlation-hierarchy).

Direction and magnitude of bias. For two binary sibling nodes X_1, X_2 with marginal success probabilities θ_1, θ_2 and covariance $\rho = \text{Cov}(X_1, X_2) > 0$ (positive correlation from a latent common cause):

Table 8.3

Node type	True probability	Independence estimate	Bias
AND	$\theta_1\theta_2 + \rho$	$\theta_1\theta_2$	$+\rho$ (underestimates — conservative)
OR	$[1 - (1 - \theta_1)(1 - \theta_2)] - \rho$	$1 - (1 - \theta_1)(1 - \theta_2)$	$-\rho$ (overestimates — optimistic)

The independence model error is exactly $\pm \text{Cov}(X_1, X_2)$, with sign determined by node type. AND-nodes are conservative because clustering of successes helps joint success. OR-nodes are optimistic because clustering of failures undermines the redundancy that OR-structure is supposed to provide. Same magnitude, opposite signs.

For strategies with OR-structure near the root (the typical case — multiple alternative paths to the objective), the net bias is overestimation: the agent believes it has more redundancy than it actually does. For AND-heavy strategies (all components must work, no alternatives), the net bias is underestimation. Mixed strategies depend on the graph topology, but OR-dominated roots are far more common than AND-dominated roots in practice.

Four regimes, in order of increasing realism:

Table 8.4

Level	Model	When correct	$\hat{P}_\Sigma$ status	Sector transfer status	Computation
L0: Independence	AND/OR propagation as-is	Causally sufficient DAG (all common causes explicit)	Correct probability	Prop B.5 (linear), B.5b (componentwise nonlinear): $\alpha_s = \alpha_c$	$O(V + E)$
L1: Augmented DAG (strict prerequisites)	Strict-prerequisite common-cause nodes added explicitly; AND/OR propagation on augmented graph	Augmented DAG is causally sufficient <i>and</i> every modeled common cause has $\theta_{\text{child} \neg C} \approx 0$	Correct for augmented graph	Prop B.6: $\alpha_\Sigma = \min(1/(n_C + 1), \theta_C \pi_{j C}/(n_j + 1))$ — three-way gating	$O(V' + E')$, larger graph
L1': Mixture form (soft facilitators)	Conditional sub-DAGs weighted by common-cause state: $\hat{P}_\Sigma = \theta_C P_\Sigma(G C) + (1 - \theta_C) P_\Sigma(G \neg C)$	Soft-facilitator common causes ($\theta_{\text{child} \neg C} > 0$) with C observable per trial	Correct for the weighted combination	Prop B.7: $\alpha_{L1'} = \min(1/(n_C + 1), \theta_C \pi_{j C}/(n_{j C} + 1), (1 - \theta_C) \pi_{j \neg C}/(n_{j \neg C} + 1))$ — five-way gating. Refuted when C unobservable (Fisher rank-1; falls back to plan-level tracking or multi-child joint observation)	$O(V' + E')$ per common cause; doubles parametric footprint for affected edges
L2: Full correlation	Arbitrary joint failure distribution over edges	Always (but requires specifying the full joint)	Correct	Reduces to L0 on the augmented joint state	Exponential in general

L0 (Independence) is the tractable baseline. All formal results in AAT's strategy layer — the sector condition transfer (Prop B.5 in R), the persistence schema (9.10), the gradient-based credit assignment (9.6) — are proved under L0. The strategy-plan-confidence error $\delta_s = \hat{P}_\Sigma - \Phi$ tracks calibration *within* the independence model; $\Phi = P_\Sigma(\theta)$ is the AND/OR formula at true edge rates, not actual plan success probability.

L1 (Augmented DAG) is the practical sweet spot — *for strict-prerequisite common causes*. The agent models correlation structure explicitly by adding common-cause nodes to Σ_t and restructuring the DAG so that the common cause is **factored above the correlation it creates**. The construction principle: place the common-cause node as an AND-prerequisite *above* the OR/AND structure whose children it correlates. This ensures that, conditional on the common cause being satisfied, the children are independent and standard AND/OR propagation is correct.

Scope of the AND-prerequisite construction: strict prerequisites only. The factoring-above principle requires that the common cause be a *strict* prerequisite — one for which $\theta_{\text{child}|\neg C} \approx 0$. Shared infrastructure going down, a required resource being absent, a gating permission being denied: in all of these, the correlated children fail when the common cause fails, so the AND-prerequisite correctly encodes the correlation. When the common cause is instead a *soft facilitator* — favorable market conditions, a supportive team culture, an enabling technology — children have $\theta_{\text{child}|\neg C} > 0$ and can succeed when the common cause is absent, just less reliably. The AND-prerequisite construction mathematically forces $P(\text{sub-plan} | \neg C) = 0$, which strictly understates success probability when C is absent. Soft facilitators therefore fall outside L1's single-pass construction and require one of:

- **Mixture form (L1')**: split the sub-plan into two conditional structures and weight by $P(C)$: $\hat{P}_\Sigma = \theta_C \cdot P_\Sigma(G | C) + (1 - \theta_C) \cdot P_\Sigma(G | \neg C)$ This keeps propagation polynomial but requires maintaining two parallel sub-DAGs

per soft-facilitator node. Per-edge credences split into two regimes ($p_{ij|C}$ and $p_{ij|\neg C}$), doubling the parametric footprint for affected edges.

- **Explicit conditioning (L2 subcase):** $R_2(G) = \sum_c P(C = c) \cdot R_2(G | C = c)$, summing over common-cause states. For k soft facilitators with binary states, this costs $O(2^k)$ in the number of common-cause states — the exponential blowup L2 was defined to name.

The gap between L1 (strict) and L2 (arbitrary) was not previously named. L1' (mixture form) fills this gap at linear cost per soft-facilitator node. The “L1 is the practical default” claim therefore applies most cleanly to strict prerequisites; for mixed environments (strict and soft common causes simultaneously), the correct construction mixes L1 factoring for strict prerequisites with L1' mixtures for soft facilitators. Treating all common causes as strict prerequisites systematically undervalues sub-plans that face soft facilitators (the opposite failure mode from L0's overestimation).

Example — strict prerequisite (see #example-L1 for full treatment): Two OR-alternatives sharing infrastructure dependency ($\theta_C = 0.8$, $\theta_{1|C} = 0.9$, $\theta_{2|C} = 0.7$, and implicitly $\theta_{1|\neg C} = \theta_{2|\neg C} = 0$ — infrastructure is a strict prerequisite). L0 computes $\hat{R}_\Sigma = 0.877$; actual is 0.776 (overestimation = ρ , the covariance from shared infrastructure). Naive L1 (common cause as parent of both alternatives, alternatives remain OR-siblings) gives the *same* overestimate because the OR-propagation still treats siblings as marginally independent. Correct L1 ($G = \text{AND}(C, G_{\text{sub}})$ where $G_{\text{sub}} = \text{OR}(A_1, A_2)$) gives the exact answer because A_1 and A_2 are conditionally independent given C and because $\theta_{1|\neg C} = 0$, so the AND-prerequisite is the correct encoding. The sector condition is verified (Prop B.6 in R) with $\alpha_\Sigma = \min(1/(n_C+1), \theta_C(1-\epsilon)/(n_{A_1}+1), \theta_C\epsilon/(n_{A_2}+1))$ — combining evidence-starvation and exploration-gating effects. B.5b transfers losslessly.

All L0 formal results transfer to correctly constructed L1 DAGs for *strict-prerequisite common causes* because the augmented DAG is a standard AND/OR DAG that satisfies causal sufficiency. For L1' (mixture form) with **observable common cause**, the formal results transfer through Prop B.7 (R) — five-way gating with $\alpha_{L1'} = \min(1/(n_C + 1), \theta_C\pi_{j|C}/(n_{j|C} + 1), (1 - \theta_C)\pi_{j|\neg C}/(n_{j|\neg C} + 1))$. For L1' with **unobservable common cause**, the per-conditional decomposition is identifiability-obstructed by the Cramér-Rao floor — the per-trial Fisher matrix is rank 1 rather than rank $2K + 1$, so no unbiased online estimator on the joint conditional vector admits a sector parameter $\alpha > 0$. This is not “open” but structurally refuted; the agent must augment C -observability, run multi-child joint observations, or fall back to L0-on-marginals (losing per-conditional diagnostics). The practical challenges are: (1) identifying which common causes matter enough to model explicitly, (2) classifying each identified common cause as strict or soft, (3) verifying C -observability for soft cases, (4) restructuring the DAG correctly for each. All four are modeling judgments, not mechanical procedures.

L2 (Full correlation) is a mathematical ideal. Specifying the full joint failure distribution over m edges requires $O(2^m)$ parameters, which violates the bounded-cognition constraint that motivates the DAG representation in the first place. L2 is useful as a reference point for characterizing the L0/L1 approximation error, not as a practical representation.

Choosing among L1, L1', and L2 requires classifying each common cause and verifying observability. $\theta_{\text{child}|\neg C} \approx 0 \rightarrow \text{L1}$ (factor above the correlation, B.6). $\theta_{\text{child}|\neg C} > 0$ with C **observable** $\rightarrow \text{L1'}$ (B.7, five-way gating). $\theta_{\text{child}|\neg C} > 0$ with C **unobservable** $\rightarrow \text{L1'}$ is identifiability-obstructed by the Cramér-Rao floor; either augment C -observability (preferred), use L2 explicit conditioning, or fall back to L0-on-marginals (losing per-conditional diagnostics). Mixed-classification environments (some strict, some soft, varying observability) combine L1 factoring for the strict and L1' mixtures for the soft observable. Treating all common causes as strict prerequisites under L1 alone systematically *understates* success on soft-facilitator branches (the symmetric failure mode to L0's overestimation); treating L0 as the default leads to systematic overconfidence — the agent believes it has more redundancy than it actually does (OR-nodes) or more fragility than it actually does (AND-nodes). L0 remains appropriate for domains with genuinely independent risks (independent parallel experiments, diversified portfolios with low correlation) and as a computational stepping stone during strategy construction.

Detecting latent common causes. An agent at L0 can detect causal insufficiency *under joint sibling observability generated by its own interventions* — not from on-policy traces alone. On-policy execution alone cannot distinguish the

with-and-without-latent-common-cause worlds (the no-go in 9.2); the unique broadly-available escape is joint sibling observability under exploration, where the agent’s interventions (7.3) generate the joint outcome data the test requires. Persistent overestimation of plan success after edge credences have converged is the *signal* of insufficiency (4.8 applied to the strategy layer); pairwise covariance among observed siblings is the *test* (positive covariance rejects independence and identifies where to add L1 nodes). Full treatment: 9.2; numerical instantiation: #example-L1. Summary references in dependent segments that elide the on-policy / interventional distinction should be treated with caution — the load-bearing scope is “joint sibling observability under exploration,” not “from observational data alone.”

Scope of the terminal construction. Terminal conditions as Boolean predicates with AND/OR aggregation work naturally for threshold, constraint, and composite objectives. For continuous-valued objectives without natural thresholds, the agent must set an operational threshold — introducing a discretization that is a practical proxy, not a lossless encoding of V_O . The primary $O_t \leftrightarrow$ theory interface remains V_O through the value object (6.6); terminal conditions are Σ_t ’s internal operational encoding.

Single-parameter edges. Each edge carries one number (p_{ij}), not two. An earlier formalism attempt used (p_{ij}, θ_{ij}) where θ was “contribution magnitude.” This was dropped because the AND/OR combination rules at nodes absorb θ ’s role — the complexity budget goes to one bit per node (γ) instead of one float per edge.

Edge-credence presentation coordinates. Each edge’s single-number credence has two equivalent presentations: the *probability-space* coordinate $p_{ij} \in [0, 1]$ (convenient for AND/OR propagation and interpretation) and the *log-odds* coordinate $\lambda_{ij} = \log(p_{ij}/(1 - p_{ij})) \in \mathbb{R}$ (the unique additive-evidence coordinate forced by the evidential-additivity axiom; see U). The two coordinates are related by the sigmoid $p_{ij} = \sigma(\lambda_{ij})$. AAT’s AND/OR status propagation and Correlation Hierarchy algebra operate in probability space; the continuous-gradient update machinery in 9.6 operates natively in log-odds and projects back to $[0, 1]$ via sigmoid at the readout interface. Props B.1–B.7 of R are stated in moment-parameter (probability-space) form; their sector-parameter content is Fisher-equivalent in either coordinate.

Acyclicity is Derived.

Each node in Σ_t represents a future event or state with temporal position $\tau_i > t$. An edge $X_i \rightarrow X_j$ requires $\tau_i < \tau_j$ (1.7: causes precede effects). A cycle $X_i \rightarrow X_j \rightarrow \dots \rightarrow X_i$ would require $\tau_i < \tau_j < \dots < \tau_i$, which is impossible for a real-valued time index.

Derived (from causal-structure + finite planning horizon).

Strategies involving iteration (“try A, if fail try B, if fail try A again”) are acyclic when time-indexed. The sequence unfolds as:

$$A_1 \rightarrow \text{check}_1 \rightarrow B_1 \rightarrow \text{check}_2 \rightarrow A_2 \rightarrow \dots \quad (8.13)$$

Each attempt is a distinct node at a distinct time. The apparent cycle is a linear chain in the unrolled view.

Formally: a finite set with a strict partial order (future events ordered by time) is representable as a DAG. This is a standard result in order theory.

Scope of the acyclicity result. This applies to Σ_t (the agent’s strategy over the future), not to M_t ’s model of the environment, which may include cyclic causal processes (feedback loops in the physical world, market dynamics, ecosystem interactions). The acyclicity is specific to the purposeful substate.

Epistemic Status.

Conditional on the 8.2 restriction. The DAG structure itself follows from operational postulates: temporal ordering (acyclicity — exact), probabilistic uncertainty (Cox’s theorem — exact; Cox 1946, “Probability, Frequency and Reasonable Expectation,” *American Journal of Physics* 14(1):1–13; Jaynes 2003, *Probability Theory: The Logic of Science*, Cambridge University Press, Chapter 2), and the Causal Markov Condition theorem under causal sufficiency (Markov factorization — proved conditional). The full argument is in Q: $P1 + P2 + \text{causal sufficiency} \rightarrow \text{CMC} \rightarrow \text{DAG}$ with Markov property. The result is *conditional on causal sufficiency* of the strategy (no latent common causes among strategy nodes). Causal sufficiency is a modeling ideal — the agent designed the graph, so *intended* causal relationships are explicit, but environmental common causes routinely go unmodeled in complex domains. When causal sufficiency fails, the Markov factorization breaks down, edge outcomes become correlated, and $\hat{R}_2$ overestimates success. This is model inadequacy (4.8), repairable by adding the missing common-cause nodes (L1 augmentation). See the Correlation Hierarchy above for the practical framework.

Causal sufficiency and edge independence. The AND/OR status propagation is correct when the DAG is causally sufficient (all common causes represented as nodes) — this follows from the Causal Markov Condition theorem (Q). In complex real-world systems, causal sufficiency is systematically violated: shared infrastructure, common-mode risks, supply chain dependencies, and correlated adversary actions introduce latent common causes. Correlated failure is the dominant case, not the exception. The Correlation Hierarchy (above) characterizes three levels: independence (L0, tractable baseline), augmented DAG (L1, practical sweet spot), and full correlation (L2, mathematical ideal). AAT’s formal strategy-layer results are proved under L0 but transfer to L1 (which is just a larger L0-compliant DAG). The independence model remains the tractable foundation for formal analysis; L1 augmentation is the recommended practice for deployment.

The AND/OR parameterization is a parsimony-motivated formulation choice within the strongly motivated graphical structure, not a derived necessity (8.2). The single-parameter edge convention is similarly a formulation choice motivated by convergence across three independent attempts.

Discussion.

The graph structure is sufficient under causal sufficiency; the parameterization is chosen. The DAG structure follows from temporal ordering (acyclicity — proved) and the Causal Markov Condition theorem under causal sufficiency (Markov factorization — proved conditional; see Q). This is a theorem-backed result, not a sketch: the CMC (Spirtes et al. 2000, Pearl 2009) proves that any causally sufficient causal DAG with independent exogenous noise satisfies the Markov factorization. The status is therefore: DAG-with-Markov-property is *sufficient* given operational postulates plus causal sufficiency — the desiderata are satisfied by this representation, but the necessity direction (whether some non-DAG structure could also satisfy them) is an open stronger result (Q). AAT uses AND/OR parameterization within this strongly motivated structure because (a) AND/OR is the most parsimonious complete basis for binary combination, (b) the representation converged across three independent formalism attempts. Alternative parameterizations within the strongly motivated graphical structure are legitimate research directions.

Combination assignment is principled but fallible. The question “if I remove one parent, can v still be achieved?” is derivable from M_t ’s causal model — it’s a principled assignment, not arbitrary. But the assignment can be wrong (false AND = pessimistic over-investment; false OR = optimistic under-investment), and should be updateable when evidence reveals a different structural relationship.

Connection to Pearl’s framework. Under causal sufficiency, the strategy DAG is a causal Bayesian network in the sense of structural causal models, not merely a structural analog: Q proves that the operational postulates (P1, P2, P4) plus causal sufficiency yield the Markov factorization via the Causal Markov Condition theorem (Spirtes et al. 2000, Pearl 2009). Pearl’s do-calculus therefore applies to status propagation and plan evaluation, with its causal content scoped by the identification regime of the data feeding the edge credences (9.5). In Regime A domains, where the agent performs genuine interventions and observes isolated outcomes, the DAG’s edge credences approximate interventional

probabilities and do-calculus yields clean causal estimates. In Regimes B and C, the edge credences are the agent’s working causal beliefs, updated from data of weaker identification strength. The DAG remains useful for planning — the agent must act on *some* causal model — but the strategy-plan-confidence score $\hat{P}_\Sigma$ inherits the identification weaknesses of its constituent edges. An agent operating primarily in Regime C should treat $\hat{P}_\Sigma$ as a rough heuristic, not a calibrated probability. When causal sufficiency itself fails (latent common causes among strategy nodes), the Markov factorization is violated and the DAG is the agent’s *intended* but uncalibrated causal model; the L1/L1’ augmentation in the Correlation Hierarchy is the principled repair.

Depth penalties on calibration. Beyond the confidence decay that deeper DAGs suffer (8.1), the two-edge strategic dynamics analysis (R (Props B.2-B.3)) shows that deeper edges are also harder to calibrate. Edge k in a chain is tested only when all upstream edges succeed, so its effective correction rate is attenuated by $\prod_{j < k} \theta_j$ (the evidence-starvation effect). Deeper DAGs therefore face a double penalty: lower aggregate confidence AND slower convergence of edge credences toward truth. This reinforces the structural pressure toward shallow, observable strategies — deep plans require both high per-edge reliability and sustained observability at every intermediate level to remain calibratable.

Edge independence is a consequence of causal sufficiency, not a separate assumption. The Correlation Hierarchy (Formal Expression, above) establishes this precisely: the AND/OR propagation’s correctness is not a matter of “assuming independence” — it is a consequence of whether the DAG is causally sufficient. The CMC theorem (Q) proves that causal sufficiency $\rightarrow$ exogenous independence $\rightarrow$ Markov factorization $\rightarrow$ correct AND/OR propagation. The assumption is causal sufficiency; independence is the consequence.

What L0 buys: Tractable $O(|V| + |E|)$ status propagation; single-parameter edges; clean persistence proofs (the sector condition transfers from per-edge credence to plan value via the Jacobian — Prop B.5 in R).

What L0 costs when the DAG is causally insufficient: $\hat{P}_\Sigma$ systematically overestimates success. The strategy-plan-confidence error $\delta_s = \hat{P}_\Sigma - \Phi$ proved persistent by B.5 tracks calibration *within the L0 model* — Φ is the AND/OR formula at true edge rates, not actual plan success (9.1). Gradient-based credit assignment (9.6) inherits the same bias: the residual $(y_G - \hat{P}_\Sigma)$ conflates per-edge miscalibration with omitted correlation structure.

The L1 remedy: Add common-cause nodes and restructure the DAG so each common cause is factored above the correlation it creates (see Correlation Hierarchy above). The AND/OR propagation then applies correctly to the restructured DAG, and all L0 formal results transfer because the restructured DAG is itself L0-compliant. The key engineering challenge is twofold: identifying which common causes matter, and positioning them correctly in the graph topology.

Relationship to Moore machines / finite-state automata — behavioral surface vs epistemic interior. The strategy DAG and the Moore machine $(S, A, \tilde{A}, \delta, \lambda, s_0)$ formalize different aspects of what “strategy” means; they are not competing representations of the same object but representations of *different* objects that both get called “strategy” in different literatures. The Moore machine encodes a *reactive policy* — given the current state and the other agent’s last action, produce an output and transition; it captures a complete input-output mapping over an indefinite horizon. The strategy DAG encodes a *causal plan* — a theory of which conditions must hold (AND/OR) and with what credence for the objective to be achieved; it captures the agent’s beliefs about the causal structure of its problem. **The Moore machine is the behavioral surface; the strategy DAG is the epistemic interior.**

Moore machine to DAG (partial embedding). A Moore machine with n states unrolls over a finite horizon H into a tree of depth H with branching factor $|\tilde{A}|$. Each path is a sequence of (state, output) pairs — a scenario. The unrolled tree is a DAG (acyclic by construction, per #deriv-graph-structure-uniqueness), but it lacks credence weights and AND/OR semantics. To produce a strategy DAG, two pieces of external information must be supplied: (a) probabilities over opponent actions at each branch (from M_t); (b) AND/OR combination rules reflecting which branches the agent treats as jointly required vs. alternatively sufficient. The Moore machine does not contain this information — it is agnostic about which opponent responses are likely or which paths matter. The Moore-to-DAG map is *injective on skeletons* but requires external information to complete.

DAG to Moore machine (lossy compilation). A strategy DAG can be compiled into a reactive policy by enumerating leaves in topological order and defining a Moore machine state per action leaf that transitions on observed success/failure. This compilation *discards* three pieces of structure: (a) the credence weights (the compiled machine is deterministic); (b) the AND/OR structure (the machine commits to a fixed execution order); (c) the causal semantics (the machine has no representation of *why* actions are ordered this way). The compiled machine executes the plan but cannot reason about it — it cannot recompute plan-confidence when a leaf credence changes, which is exactly what status propagation (`#disc-credit-assignment-boundary`) provides for the strategy DAG.

Neither representation is strictly more general; the two are *orthogonal*. Moore machines can represent reactive strategies with no causal model (tit-for-tat, grim trigger) that have no natural DAG form — there is no plan to encode, just a response rule. Strategy DAGs can represent plans with rich conditional structure and uncertainty that would require exponentially many Moore-machine states (because the Moore machine must enumerate every possible observation sequence, while the DAG factors through conditional independence). The orthogonality is structural: behavioral surface vs epistemic interior. The DAG’s acyclicity is not a restriction — it is a consequence of temporal indexing per `#deriv-graph-structure-uniqueness`; what looks like a Moore-machine cycle (revisiting state “cooperate” at $t = 1$ and $t = 5$) is two distinct nodes $v_{t=1}, v_{t=5}$ in the time-unrolled DAG. For finite horizons the two have equivalent expressive power over behavioral sequences, with the Moore machine exponentially more compact for repetitive strategies and the DAG exponentially more compact for plans factoring through conditional independence.

Composition behaviour. Miller’s product-automaton construction gives *exact* composition: two Moore machines interacting produce a meta-machine with state space $S_1 \times S_2$ and deterministic transitions. AAT’s composition closure (`#form-composition-closure`) uses an *approximate* dynamical homomorphism with closure defect ϵ^* . These address different questions. The product automaton asks “what behaviour does the pair produce?” — answer: another automaton, exactly. AAT asks “can this pair be described as a single AAT agent?” — answer: approximately, with bounded error. Even if strategies were Moore machines, composition of *behaviour* would become exact (product automaton) but composition of *agent descriptions* (M_t, O_t, Σ_t) would still require the approximate framework — the closure defect comes from projecting the joint internal state, not from composing the policy.

Why AAT uses the DAG rather than the Moore machine. The defining feature of an adaptive agent (`#scope-adaptive-system`) is that the agent revises its strategy when M_t changes. Strategy revision requires the agent to know *why* it is doing something, not just *what*. A Moore machine specifies what to do but contains no theory of why this is the right thing to do — when M_t changes, the Moore machine has no internal handle to revise from. The strategy DAG carries the causal semantics that strategy revision operates on (per `#disc-credit-assignment-boundary` and the gradient-based candidate signal); compiling to a Moore machine would discard the very structure adaptation requires.

Composing heterogeneous strategy DAGs (causal-abstraction typing). The composition-behaviour distinction above — that composing *agent descriptions* needs the approximate framework while composing *behaviour* does not — has a precise mechanism at the strategy layer, and it is order-theoretic, not a magnitude in team size. Composing strategy DAGs $\Sigma_1, \dots, \Sigma_N$ into one macro-strategy Σ_c is the causal-model-abstraction question with the arrow reversed: not “when does one high-level model abstract one low-level model?” but “when does one high-level strategy DAG Σ_c faithfully abstract the *joint* micro-strategy $\Sigma_1 \times \dots \times \Sigma_N$?” Since the strategy DAG is a causal Bayesian network under causal sufficiency (the “Connection to Pearl’s framework” Discussion above), the right object is a *strategy τ -abstraction* — a Σ_c with a surjective node map τ and an intervention map under which the status-propagation diagram commutes — adopting the causal-abstraction framework of Rubenstein et al. (2017), Beckers & Halpern (2019), and Beckers, Eberhardt & Halpern (2020). The composition-closure defect ϵ_Σ of `#form-composition-closure` and the Beckers–Eberhardt–Halpern approximate-abstraction error are the *same commuting-diagram family but not term-for-term*: the BEH error is a supremum over all interventions and contexts (interventional), ϵ_Σ is an expectation over the reachable on-policy trajectory distribution (observational), so a (τ, α) -abstraction implies a bounded ϵ_Σ but not conversely, and BEH leaves its high-level distance application-dependent — the norm-on-graphs choice remains AAT’s. The adoption buys the correct type and a strictly-stronger optional target, not a closed defect formula.

Exact composition is an order-compatibility condition, not a small- N regime. An exact Σ_c ($\varepsilon_\Sigma = 0$) exists *iff* the joint micro-strategy's interventional distribution factors through τ — equivalently, *iff* on shared nodes the sub-plans' temporal sub-orders are jointly acyclic *and* their AND/OR root semantics agree (the root combinator is dictated by the composite objective via #scope-composite-agent, not free). When these hold, Σ_c is exact for *all* N : compatible heterogeneous strategy composition is dimension-free, the strategy-layer twin of the dimension-free state-composition regime. The agent count never enters. The complementary fixed-topology case — agents sharing one skeleton and disagreeing only on edge-credences — has its own exact closed-form defect ($\varepsilon_\Sigma^{*2} = |E| \cdot \overline{\text{Var}_i[\eta_{\Sigma,i}]} \cdot \text{Var}[r]$, also N -free), landed as the strategy-layer instance of the structural-unity axis in #result-unity-closure-mapping.

When shared sub-orders conflict, no exact single-DAG macro-strategy exists. If shared nodes S carry incompatible sub-orders (agent 1 plans $A < B$, agent 2 plans $B < A$ on shared A, B), the union of the per-agent partial orders is cyclic and this segment's derived-exact acyclicity ("Acyclicity is Derived", above) *forbids* an exact Σ_c — non-existence, not a large defect. The minimal admissible Σ_c is the **condensation** of the union-order graph by strongly connected components (the unique minimal acyclic quotient; Tarjan): each cyclically-entangled set of shared nodes collapses to one coarse macro-node the composite cannot temporally distinguish. The information destroyed by collapsing an SCC C is exactly the total correlation among the collapsed status variables, so the irreducible composition defect is

$$\varepsilon_\Sigma^{*2} = \sum_{C \in \text{SCC}(G_U), |C| > 1} \text{TC}(\{s_v\}_{v \in C}) \leq |S| \log 2, \quad (8.14)$$

zero *iff* the collapsed statuses were independent (the order was spurious), maximal for a deterministic chain (the order was real) — the strategy-layer analog of the Mori–Zwanzig projection residual that governs state composition, with the SCC-condensation as the projection. The bound is **shared-plan-size-bounded and fully N -free**: redundant order-conflicts among many agents collapse to the same SCCs (the condensation is idempotent under implied constraints); only *fresh* shared nodes per agent grow the defect, and that growth is linear in plan overlap and attributable to plan design, never to the composition operator. When even the condensation over-coarsens past usefulness, the honest macro-object is a mixture over per-context conditional sub-DAGs — structurally the L1' mixture form of the Correlation Hierarchy above — whose support (the number of incompatible order-contexts) is itself bounded by the SCC count $\leq |S|$, not by N .

Downstream corollary (a Brooks's-Law floor). Because ε_Σ^* is $|S|$ -bounded and N -free, it cannot by itself flip the composite persistence inequality (CM2) of #deriv-critical-mass-composition at large N (it enters only through $\rho_{\text{eff}} = \rho_{\text{ext}} + \varepsilon^* \nu_c$): strategy-topology heterogeneity *alone* cannot cause a Brooks's-Law collapse — only coordination overhead can — consistent with #disc-composition-consistency's Tier-1M closed form.

Epistemic status of this subsection. The τ -abstraction typing is a prior-art adoption (high confidence on the type; the BEH-vs- ε_Σ relation is BEH $\Rightarrow$ bound, with the supremum-vs-expectation / interventional-vs-observational reconciliation named and deliberately left open). The exactness-*iff* condition and the non-existence under cyclic shared sub-orders rest on this segment's own derived-exact acyclicity. The SCC-condensation total-correlation defect law is solid in structure (the condensation is the classical unique minimal acyclic quotient; total correlation is the collapse loss by the log-odds chain-rule identity) and *conditional on its constant*: whether the $|S| \log 2$ ceiling is tight when intra-SCC statuses are strongly but non-deterministically coupled is one named open sub-question (Working Notes). The subsection is therefore *conditional*, matching the #form-composition-closure row it makes concrete.

Working Notes.

- The Correlation Hierarchy (L0/L1/L2) is now first-class in the Formal Expression. The relationship between causal sufficiency, edge independence, and $\hat{R}_\Sigma$ accuracy is grounded by the CMC theorem. The main open question is practical: what heuristics help agents identify which common causes are worth modeling at L1? This is a domain-specific engineering question, not a theoretical gap.

- The graph-sufficiency argument (P1, P2, P4 + causal sufficiency $\Rightarrow$ DAG with Markov property via CMC) is proved, not sketched. See Q for the full derivation and the parallel to Cox's theorem (Cox is necessary-and-sufficient; this argument is sufficient only). The DAG-with-Markov-property is sufficient given the postulates plus causal sufficiency; the AND/OR parameterization remains a formulation choice, strategy-dag stays typed as Definition because the *parameterization* is chosen, and the necessity of the *graphical structure* itself is an open stronger result.
- Health metrics (groundedness, observability coverage, weighted redundancy, bottleneck scores) are scaffold — engineering quantities for monitoring DAG health, not principled derivations. They may be useful for implementation but should not enter the theory's formal chain.
- **Satisfaction criterion.** $V_{O_t}^{\min}$ is now introduced in 6.4 as a parameter of the objective — the minimum acceptable trajectory value. The well-formedness constraint references it from there, not from 8.4. The satisfaction gap diagnostic builds on $V_{O_t}^{\min}$ but does not define it.
- **Terminal alignment error.** When the agent achieves its terminal conditions but evaluates $V_{O_t}(\tau) < V_{O_t}^{\min}$ on the actual trajectory, the well-formedness belief was wrong — the operational success criteria didn't capture what the objective actually required. This is detectable only through experience (achieve the terminals, evaluate V_{O_t}), not through a priori analysis. It triggers terminal reassessment — a structural change in Σ_t driven by the $O_t \leftrightarrow$ terminal mismatch. Whether this should be formalized as a named diagnostic signal (δ_{align}) alongside δ_{sat} , δ_{regret} , and $\delta_{\text{strategic}}$ is open.
- **Open (composition defect-constant tightness).** The SCC-condensation defect law in the composition subsection has a clean ceiling $\varepsilon_S^{*2} \leq |S| \log 2$; the *structure* (condensation + total-correlation collapse loss) is solid, but whether the ceiling is tight when intra-SCC status variables are strongly but non-deterministically coupled is not pinned. The closing instance is a two-shared-node worked case ($S = \{A, B\}$, agent 1 $A < B$, agent 2 $B < A$, Beta-Bernoulli credences) — the strategy-layer analog of the correlated-Kalman numerical instance. Moderate effort; resolving it lifts the defect-constant from conditional to exact. (Indexed: spikes/PROPOSED.md Tier 1 — “SCC-defect-constant tightness (A1 Q1)”.)
- **Open (mixture-support dynamics).** When the SCC-condensation over-coarsens and the macro-object is the L1'-style mixture over incompatible order-contexts, its support is bounded by the SCC count $\leq |S|$ *statically*. Whether that support can *grow under the Orient cascade* (agents revising shared sub-orders as M_t changes — terminal reassessment, above) is the one place a genuine dynamical N -or-time effect could re-enter; the conjecture is that it stays SCC-count-bounded while the cascade reshuffles which nodes are shared. A follow-on, not a present claim. (Indexed: spikes/PROPOSED.md Tier 2 — “Mixture-support dynamics under the Orient cascade (A1 Q2)”.)

Incidental audit gold (lift 2026-05-31, A8 batch). Cross-audit “wandering thoughts” / §14-ideation harvested from the de-novo auditors' working dirs, deduplicated across substrates and lightly attributed. Orthogonal pedagogical / framing / figure / forward-vision material, kept separate from certified theory-fix findings. **Coverage:** this was the cluster's heaviest segment — 9 dirs reached a digested reflection (Gemini, AUDIT-WORKING-193847; Claude, AUDIT-WORKING-266847; Codex/Claude, AUDIT-WORKING-361742; Codex/Claude, AUDIT-WORKING-526815; Claude, AUDIT-WORKING-584721; Claude, AUDIT-WORKING-773921; Gemini, AUDIT-WORKING-829314; Gemini, AUDIT-WORKING-849201; Claude, AUDIT-WORKING-451729-batch-09 + 471203-batch + 963715-batch). Felt value clustered at top-decile / “10-of-10 load-bearing” across substrates.

1. Candidate Brief prose / pre-prose.

- **The AND/OR bias as one symmetric sentence.** “AND-nodes are *conservative* (independence underestimates joint success by $+\rho$ — clustering of successes helps); OR-nodes are *optimistic* (independence overestimates by $-\rho$ — clustering of failures defeats the very redundancy OR is supposed to provide). Same magnitude, opposite signs.” Multiple substrates re-derived and praised this as the segment's crispest takeaway (Gemini, AUDIT-WORKING-193847; Claude, AUDIT-WORKING-266847; verified algebraically at Claude, AUDIT-WORKING-451729 / 471203 / 584721). The body already states it; this is the Brief-grade compression.
- **“ $\hat{P}_2$ is a heuristic, almost always optimistic” as the one-line caution.** “Any proof relying on $\hat{P}_2$ being perfectly calibrated is suspect unless it explicitly claims L1/L2 causal sufficiency” (Gemini, AUDIT-WORKING-193847).

2. Candidate Discussion.

- **Project-management overconfidence as the canonical Correlation-Hierarchy story.** The single most-converged Discussion candidate (3+ substrates). “Project managers use L0 planning: three independent backup teams at 70% each, so $1 - (1 - 0.7)^3 = 97.3\%$. But the three teams share a staging server / AWS region / one manager’s attention — a latent common cause C . Positive correlation makes L0 *overestimate* the OR-gate by exactly ρ ; true success might be $\sim 80\%$. The fix is mechanical: add the shared dependency as a node and place it as an AND-prerequisite *above* the OR-gate — $\text{Plan} = \text{AND}(\text{Staging works}, \text{OR}(\text{Team1}, \text{Team2}, \text{Team3}))$ — and the math is exact again” (Gemini, AUDIT-WORKING-829314; same mapping at Gemini, AUDIT-WORKING-193847 for “defense in depth in cybersecurity/aerospace”; Claude, AUDIT-WORKING-451729 for “business-continuity / shared-database / military shared-logistics”; Claude, AUDIT-WORKING-773921 “why project managers systematically underestimate risk”). A ready worked-example Discussion paragraph; #example-L1 already carries the numerics, so this is the *narrative* gloss that motivates them.
- **“Highly-redundant ‘safe’ architectures are more fragile than their designers believe.”** The general lesson under the bias analysis: heavy OR-gates (redundancy) are *less* robust than they look under unmodeled positive covariance, while heavy AND-chains are slightly *more* robust than feared — “when an agent or organization is surprised by a failure, it is usually because an unmodeled common cause collapsed an OR-gate it thought provided safe redundancy” (Gemini, AUDIT-WORKING-829314).
- **Moore-machine vs DAG = behavioral surface vs epistemic interior.** The existing “Relationship to Moore machines” subsection was independently singled out as “the clearest explanation of Reactive Policy vs Causal Planning I have ever encountered — a Moore machine says *what* to do; a DAG says *why*, and you cannot do strategy revision without knowing *why*” (Claude, AUDIT-WORKING-773921). A signal that this subsection is already pedagogically strong and a Brief/preamble could lean on the “what vs why” line.

3. Follow-up items.

- **Formalize the terminal-alignment error as a named diagnostic δ_{align} .** Strong convergence (4 substrates) on the Working-Notes “terminal alignment error” deserving promotion to a fourth named diagnostic signal alongside δ_{sat} / δ_{regret} / $\delta_{\text{strategic}}$: the case where the agent achieves all terminal conditions ($\hat{R}_T \approx 1$) yet $V_{O_t}(\tau) < V_{O_t}^{\min}$ on the realized trajectory — “the terminals were a bad proxy for O_t .” Glossed as “the math of a mid-life crisis” (Gemini, AUDIT-WORKING-193847), “King Midas / be-careful-what-you-wish-for” (Gemini, AUDIT-WORKING-829314), and flagged as open-to-formalize at Claude, AUDIT-WORKING-266847 / 361742. The body’s Working Notes already raise this as open; the cross-substrate convergence is the signal it is worth doing.
- **Relocate the Correlation Hierarchy from Formal Expression into Discussion.** Two substrates: the Hierarchy is the most important *conceptual* addition but reads as a practical-deployment framework rather than a definitional claim, and is “buried in Formal Expression after substantial structural definition material” — candidate for Discussion placement, or even its own Derived/standalone segment given its depth (“almost deserves its own segment”, Gemini, AUDIT-WORKING-829314; placement note at Claude, AUDIT-WORKING-451729 / 471203). A structure question, not a content gap.
- **Topology discovery is unaddressed.** “The segment explains how to propagate probabilities and how to repair broken topologies (L1 augmentation), but not where the initial nodes and edges come from — does the agent use an LLM prior? random path generation? Topology discovery seems like a missing piece of the causal-learning puzzle” (Gemini, AUDIT-WORKING-193847; echoed at Gemini, AUDIT-WORKING-849201 “how does the agent detect an unmodeled common cause exists?” → #der-causal-insufficiency-detection). A scope-of-DAG-acquisition follow-up.

4. Readers often ask / wonder.

- **How does a bounded agent actually compute the supremum behind $\hat{R}_T$ / A_O ?** Recurs across the cluster — finding the optimal policy is NP-hard, so the agent must maintain a running lower-bound estimate via heuristic search/deliberation rather than a true supremum (Gemini, AUDIT-WORKING-829314, raised on the sibling sat-gap segment but lands on $\hat{R}_T$ here too).

- **What about vector / Pareto / multi-terminal objectives?** The single-root + scalar-threshold construction fits scalar objectives; genuinely multi-terminal or Pareto objectives need the scalarization / AND-terminal workaround called out explicitly (Codex/Claude, AUDIT-WORKING-526815).

5. Candidate figures.

- **Split view: propagation path + correlation hierarchy, with a “bridge warning.”** The clearest diagram puts the main DAG status-propagation path on one side and the L0/L1/L1'/L2 hierarchy on the other, with an explicit annotation at the causal-sufficiency $\rightarrow$ AND/OR-propagation step warning that causal sufficiency buys the *Markov factorization*, while the specific noisy-AND/noisy-OR product form is an *additional* local parameterization choice (Codex/Claude, AUDIT-WORKING-526815). (See the off-ramp below — this annotation is the figure-form of a real precision point.)

Belongs elsewhere.

- **ELI / Section IV — “survive success.”** “Many RL agents terminally halt when they reach their goal state; but if the goal was poorly formulated, halting is death. A persistent consciousness must be able to experience the terminal-alignment error, realize its DAG was a bad proxy for its true O_t , and structurally rewrite the DAG *without deleting itself* — the ability to survive the realization that your life’s work aimed at the wrong target” (Gemini, AUDIT-WORKING-193847). Aspirational reach toward 04-eli-core/, anchored on the terminal-alignment-error concept.

Off-ramp (NOT gold — flagged for the certified track).

- **Causal sufficiency $\Rightarrow$ Markov factorization, *not* $\Rightarrow$ the noisy-AND/OR product formulas.** Two substrates independently flagged that several passages read as “AND/OR propagation is correct iff the DAG is causally sufficient, because causal sufficiency guarantees independent edge outcomes” — but CMC/causal sufficiency gives *conditional independence and a Markov factorization over variables*, not the specific per-edge independent-contribution product model. Correct propagation needs *both* causal sufficiency *and* the local AND/OR parameterization to actually hold. Suggested three-layer separation: (1) causal sufficiency $\rightarrow$ causal-Bayesian-network factorization; (2) AND/OR $\rightarrow$ a *chosen* compact local parameterization; (3) Correlation Hierarchy $\rightarrow$ repairs violations of the independence assumption *within* that parameterization (Codex/Claude, AUDIT-WORKING-526815; Claude, AUDIT-WORKING-451729 noted the same “proved vs sufficient” boundary). The chapter-end #impl-strategy-structure reportedly already states the cleaner three-layer version; the local #def-strategy-dag Epistemic-Status / “edge independence is a consequence of causal sufficiency” wording is the straggler — routed for the lead to reconcile (a register/precision fix, not a status change).
- **The blanket “ $\hat{R}_2$ systematically overestimates” claim is topology-dependent.** The bias table itself shows positive covariance makes OR-estimates optimistic but AND-estimates *conservative*; the blanket “overestimates” should be scoped to OR-dominated-root structures (which the body argues are the common case — so the fix is to state the scope, not drop the claim) (Codex/Claude, AUDIT-WORKING-526815). A scoping tightening.

8.4 Definition: Satisfaction Gap

The first of two orthogonal diagnostic quantities that together separate “the goal is too hard” from “the strategy is too weak” (8.5 is the second). The framework first defines **objective attainability** A_O as the best achievable value given the agent’s current beliefs M_t , available policy class Π , and planning horizon N_h — the supremum of the value object over policies. The **satisfaction gap** δ_{sat} is then the distance between the objective’s satisfaction threshold $V_{O_t}^{\min}$ and this attainability: positive when the objective is *unmet* under the best available policy given the current model and horizon; non-positive when the objective is attainable in principle.

A positive δ_{sat} does not automatically mean the goal is wrong. The signal has multiple possible causes — genuinely infeasible objective, too-narrow policy class, too-short horizon, model wrong about feasibility, or jointly-infeasible compound terminals — each with a different remedy. **Objective revision is the last resort, not the first response to unmet goals.** The orient cascade (10.1) formalizes this ordering, and the disambiguation table below encodes the diagnostic procedure.

Under the canonical continuation convention (6.6), δ_{sat} is a *local* diagnostic — it answers “can I improve toward the goal from here?” not “is the goal globally feasible?” The gap value is **convention-relative**: the C1 / C2 / C3 hierarchy on continuation conventions induces a monotonicity hierarchy on the gap (C1 most conservative, C3 strongest), so analyses using different conventions cannot be directly compared. See Epistemic Status.

Formal Expression.

Definition (objective-attainability).

$$A_O(M_t; \Pi, N_h) = \sup_{\pi \in \Pi} V_O(M_t, \pi; N_h) \quad (8.15)$$

The **objective attainability** — the best achievable value given current beliefs M_t , available policy class Π , and horizon N_h .

Definition (satisfaction-gap).

$$\delta_{\text{sat}} = V_{O_t}^{\min} - A_O(M_t; \Pi, N_h) \quad (8.16)$$

where $V_{O_t}^{\min}$ is the minimum trajectory value that counts as “objective met” — a threshold set by the objective itself (for constraints: all satisfied; for utility: a minimum acceptable level).

- $\delta_{\text{sat}} > 0$: The objective is **unmet** under the best available policy, current model, and horizon.
- $\delta_{\text{sat}} \leq 0$: The objective is **attainable** in principle.

Disambiguation. $\delta_{\text{sat}} > 0$ does NOT automatically mean the goal is wrong. It means the goal is unmet given (M_t, Π, N_h) . The positive signal has multiple possible causes:

Table 8.5

<i>Cause</i>	<i>Fix</i>	<i>How to distinguish</i>
Goal is genuinely infeasible	Revise O_t	Persists across M_t, Π, N_h improvements
Policy class too narrow	Expand Π (structural adaptation of Σ_t)	δ_{sat} decreases when richer policies are tried
Horizon too short	Extend N_h	δ_{sat} decreases with longer planning horizon
Model is wrong about feasibility	Improve M_t (reduce $\delta_{\text{epistemic}}$)	δ_{sat} changes when M_t is corrected
Objectives jointly infeasible	Revise O_t or relax constraints	Individual terminal satisfaction gaps are zero but AND-node fails; cross-terminal tradeoff is required

Objective revision is the **last resort**, not the first response to unmet goals. The orient cascade (10.1) formalizes this ordering.

Epistemic Status.

Exact as a definition — convention-relative as a diagnostic. The satisfaction gap is a mathematical definition — the difference between a threshold and a supremum over a function class. The definition is precise; the *computation* of A_O is generally intractable (it requires optimization over a policy class), but the quantity is well-defined. However, the *value* of δ_{sat} depends on three external parameters: the continuation convention (π_{cont}), the horizon (N_h), and the scalarization of the objective (V_{O_t}). These are not derived by AAT — they are choices the analyst makes. The satisfaction gap is therefore an intrinsic architectural diagnostic *given* a measurement convention, not an absolute property of the agent.

Convention dependence and the hierarchy. A_O inherits the continuation convention from the value object (6.6), which defines three named conventions forming a monotonicity hierarchy:

- **C1** (one-step improvement, canonical): $\delta_{\text{sat}}^{(1)} = V_{O_t}^{\min} - A_O^{(1)}$. Tests whether the agent can improve toward the goal in one step. Most conservative — a multi-step recoverable objective may show $\delta_{\text{sat}}^{(1)} > 0$ because continuation is frozen at π_{current} .
- **C2** (receding-horizon): $\delta_{\text{sat}}^{\text{RH}} = V_{O_t}^{\min} - A_O^{\text{RH}}$. Tests whether the agent can reach the goal with N_t -step replanning. Captures recovery paths invisible to C1.
- **C3** (Bellman): $\delta_{\text{sat}}^{\text{B}} = V_{O_t}^{\min} - A_O^{\text{B}}$. Tests whether the goal is genuinely infeasible given (M_t, Π, N_h) . Strongest diagnostic — $\delta_{\text{sat}}^{\text{B}} > 0$ means no policy in Π can achieve the objective.

The monotonicity result (6.6): $\delta_{\text{sat}}^{\text{B}} \leq \delta_{\text{sat}}^{\text{RH}} \leq \delta_{\text{sat}}^{(1)}$. The $\delta_{\text{sat}}^{\text{B}} \leq \delta_{\text{sat}}^{\text{RH}}$ half is unconditional; the $\delta_{\text{sat}}^{\text{RH}} \leq \delta_{\text{sat}}^{(1)}$ half holds under the order-consistency condition on C2's replanning objective (RH-1/2/3 in 6.6), and can fail for unguarded short-horizon replanning. C1 gives the most false “unattainable” diagnoses; C3 gives none (modulo model error in M_t). Analyses using different conventions should not be compared directly — the convention is part of the measurement.

Discussion.

Why two gap measures, not one. An earlier formulation used a single $\delta_{\text{objective}}$ for goal-related mismatch. This conflates two distinct situations: “the goal is too hard” and “the strategy is too weak.” When the agent is optimally pursuing an infeasible goal, $\delta_{\text{objective}}$ is large but there's no strategy to improve — the problem is the goal, not the plan. The satisfaction gap (δ_{sat}) and control regret (8.5) separate these cases, enabling the right corrective action. This is the **anti-collapse** discipline (5.7) again — the same move met at β vs ρ (2.2), now in its two-quantities form: a single goal-distance signal would merge two situations that route to *opposite* repairs (objective revision vs strategy revision), so the split is what makes the orient cascade's response correct rather than a guess.

The disambiguation table is load-bearing. Without it, an agent facing $\delta_{\text{sat}} > 0$ might immediately revise its objective when the real problem is an inadequate model or a too-narrow policy class. The table encodes the diagnostic procedure: check M_t adequacy first (maybe the goal IS feasible but the model doesn't know it), then check Π and N_h , and only then consider revising O_t .

Productive infeasibility — the aspiration hypothesis. [*Hypothesis (aspiration-dominance)*] A persistent $\delta_{\text{sat}} > 0$ is not always a defect to be cascaded away: maintaining an infeasible aspirational objective O_a can be optimal *by a feasible objective's own measure*. Writing π_a and π_f for the optimal policies under aspirational O_a ($\delta_{\text{sat}} > 0$) and feasible O_f ($\delta_{\text{sat}} = 0$), the claim is that under nameable conditions

$$V_{O_f}(M_t, \pi_a) > V_{O_f}(M_t, \pi_f) \quad (8.17)$$

— pursuing the unreachable target outperforms direct pursuit of the reachable one. Candidate mechanisms (each a known phenomenon in optimization practice, none yet derived here): the aspirational gradient cuts through local

optima of O_f 's landscape; the aspirational objective implicitly regularizes the policy space against shortcut policies; and the aspirational gradient can be better-conditioned near π_f than O_f 's own. If this holds, the cascade's rule sharpens: objective revision is warranted not merely when $\delta_{\text{sat}} > 0$ persists across $M_t/\Pi/N_h$ improvements (the table's last row), but only when the aspirational gradient is *worse* than the best feasible gradient — “shoot for the stars, land on the moon” as a claim about trajectory quality through strategy space, not motivational rhetoric. Hypothesis-grade: the inequality is stated, the conditions are informal, and the connection to optimistic exploration / over-parameterization effects in the learning literature needs a relata pass. Structure-modifying version (aspirational *models*, not just objectives): 7.8.

Dependence on M_t . A_O is computed from M_t , not from the true environment state Ω_t . The agent's assessment of attainability could be wrong — an achievable goal might look unachievable with a bad model, or vice versa. Improving M_t (reducing $\delta_{\text{epistemic}}$) brings the agent's attainability assessment closer to reality. This is why the orient cascade puts epistemic update before attainability evaluation.

Diagnostic content vs. AI's expected-free-energy decomposition. Active inference's expected free energy (EFE) decomposes into *pragmatic value* (how preferred are the outcomes the policy expects?) and *epistemic value* (how much does the policy reduce uncertainty?) (Friston, FitzGerald, Rigoli, Schwartenbeck & Pezzulo 2017, “Active inference: a process theory,” *Neural Computation* 29; Da Costa, Parr, Sajid, Veselic, Neacsu & Friston 2020, “Active inference on discrete state-spaces,” *J. Math. Psych.* 99 §2.4; Sajid, Ball, Parr & Friston 2021, “Active inference: demystified and compared,” *Neural Computation* 33). The decomposition supports policy ranking but does not separate two distinct diagnoses that AAT's apparatus does separate: “the goal is unattainable from here” ($\delta_{\text{sat}} > 0$, this segment) versus “the current policy is not the best available” ($\delta_{\text{regret}} > 0$ in 8.5). Both increase EFE without distinguishing the cause. The 2×2 cell map in the disambiguation table above gives the four diagnoses the orient cascade (10.1) acts on differently — strategy revision, objective revision, action vs. learning. AI's pragmatic-epistemic split does not produce this disambiguation.

The diagnostic structure depends on V_{O_t} being a *value functional* on trajectories (6.4) and A_O being an *attainability supremum*, not on outcomes encoded as log-priors — AI's preferences-as-priors form ($C(o) = \log P_{\text{pref}}(o)$) collapses the diagnostic by making “wanting o ” and “expecting o ” formally the same operation (the dark-room critique, Sun & Firestone 2020, “The dark room problem,” *Trends Cog. Sci.* 24). Sun & Firestone diagnose the preferences-as-priors collapse; AAT's value-functional reformulation is AAT's own downstream architectural response to that diagnosis, not a move Sun & Firestone themselves propose. The conservative reading of the citation: AAT takes their critique as motivation for abandoning the preferences-as-priors form in favor of value functionals that separate outcome value from outcome prediction, enabling the diagnostic structure the 2×2 cell map requires. This is a deliberate divergence from AI, not an oversight.

Working Notes.

- $V_{O_t}^{\text{min}}$ (the satisfaction threshold) is a property of the objective, not of the agent. For constraint-satisfaction objectives, it's natural (all constraints met = satisfied). For utility-maximizing objectives, it's less obvious — what counts as “good enough”? This threshold may need to be explicitly modeled as part of O_t .
- The supremum in A_O is over Π , the available policy class. For agents with explicit Σ_t , Π corresponds to the set of strategies representable in the agent's DAG formalism. Expanding Π (structural adaptation of Σ_t — adding nodes, edges, or changing γ assignments) is the purposeful analog of 4.8.
- **Cross-reference to NeurIPS Paper 2.** The satisfaction-gap / control-regret two-gap diagnostic is **Component 1** of NeurIPS 2026 Paper 2's composition theorem (“A Unified Convergence Theory for Non-Stationary Reinforcement Learning”, [~/src/neurips/02-unified-convergence-rl/](#), §4 Theorem 4.1). The diagnostic separates “you're not doing it well enough” (control regret) from “the world doesn't permit it” (satisfaction gap), preventing dark-room collapse in the composed regret bound. The paper's hypothesis (A1) carries the extended-real reading required for the two-gap structure to hold even when the satisfaction frontier is empty. See [spikes/neurips-back-integration-2026-05-08.md](#) §1 Paper 2.

Incidental audit gold (lift 2026-05-31, A8 batch). Cross-audit “wandering thoughts” / §14-ideation harvested from the de-novo auditors’ working dirs, deduplicated across substrates and lightly attributed. Orthogonal pedagogical / framing / figure / forward-vision material kept separate from certified theory-fix findings. **Coverage:** 9 dirs reached a digested reflection on this segment, mostly as the first half of the satisfaction-gap/control-regret pair (Gemini, AUDIT-WORKING-193847; Claude, AUDIT-WORKING-266847 — paired sat+regret note; Codex/Claude, AUDIT-WORKING-361742; Codex/Claude, AUDIT-WORKING-526815; Claude, AUDIT-WORKING-584721; Claude, AUDIT-WORKING-773921; Gemini, AUDIT-WORKING-829314; Gemini, AUDIT-WORKING-849201; Claude, AUDIT-WORKING-451729-batch-10 + 963715-batch + 471203-batch). High convergence: the disambiguation table is read across substrates as the segment’s load-bearing operational core. Substrate attribution inferred from voice where not explicit.

1. Candidate Brief prose / pre-prose.

- **“Formalizes the feeling of ‘this goal is impossible.’”** Convergent one-line gloss of what $\delta_{\text{sat}} > 0$ under the strongest convention means — the diagnostic that lets an agent recognize it is “trapped in an unwinnable game rather than endlessly optimizing a doomed strategy” (Gemini, AUDIT-WORKING-849201; Claude, AUDIT-WORKING-773921). A plain-language Brief anchor for the satisfaction-gap half of the pair.
- **The map-with-a-missing-bridge gloss of the disambiguation step.** “ $\delta_{\text{sat}} > 0$ does not mean give up — first ask: am I sure it’s impossible, or is my map just missing a bridge?” before blaming the goal (Gemini, AUDIT-WORKING-829314). A Feynman-grade re-statement of “objective revision is the last resort.”

2. Candidate Discussion.

- **The diagnostic disambiguation sequence as the mathematics of grit vs. acceptance** (*strongest convergent Discussion candidate — 3 substrates*). The order-of-operations the disambiguation table encodes — improve $M_t \rightarrow$ extend $N_h \rightarrow$ expand $\Pi \rightarrow$ revise O_t (Epistemology $\rightarrow$ Horizon $\rightarrow$ Strategy $\rightarrow$ Objective) — is read as a formalization of resilience: “an agent that revises its objective *first* is a pure satisficer with zero perseverance; an agent that loops forever trying to improve its model when the goal is genuinely thermodynamically impossible is a stubborn fool who dies of starvation while studying a brick wall. The framework provides the exact mathematical boundaries for when an agent should transition from stubbornness to acceptance” (Gemini, AUDIT-WORKING-829314). Same framing as “Grit vs. Acceptance” at Gemini, AUDIT-WORKING-193847 (“if $\delta_{\text{sat}} > 0$ under C3 the goal is impossible and grit is pathological; if positive only under C1 the goal is merely hard and grit — expanding Π — is optimal”) and the psychology read at Claude, AUDIT-WORKING-266847. The body’s Discussion already states the ordering; this is the *narrative motivation* for why the ordering is the right one.
- **The un-updated- O_t / updated- M_t split as the reason the diagnostic survives the dark room.** Why the value-functional reformulation matters, said vividly: “if your goal is encoded as a strong prior belief about the world (preferences-as-priors), then when you face harsh reality your Bayesian update naturally weakens your ‘belief’ in the goal — you stop wanting things because reality proved you don’t have them. By strictly enforcing $X_t = (M_t, G_t)$, AAT lets the agent acknowledge a terrible reality (M_t updates to ‘I am starving’) while maintaining an ironclad un-updated objective (O_t remains ‘I want to eat’). This irreducible mismatch $\delta_{\text{sat}} > 0$ is the engine that drives all subsequent action” (Gemini, AUDIT-WORKING-829314). A candidate Discussion gloss of the existing Sun & Firestone paragraph that makes the dark-room point concrete.
- **Convention-relativity is measurement-theoretic, not merely computational.** Sharpest framing of the convention-dependence caveat: “two analyses computing δ_{sat} under different conventions aren’t disagreeing; they’re measuring different things — closer to how empirical sciences treat measurement protocols than how mathematics treats theorems. AAT is willing to be measurement-theoretic about its own quantities” (Claude, AUDIT-WORKING-471203-batch; same point at Claude, AUDIT-WORKING-584721 — “the diagnostic is meaningful *within* a convention; comparing across conventions is incoherent”). The Epistemic Status already says the convention is part of the measurement; this names *why* that is a substantive stance.

3. Follow-up items.

- **The two gaps sum to total shortfall — a candidate clarifying identity.** $\delta_{\text{sat}} + \delta_{\text{regret}} = V_{O_t}^{\min} - V_O(\pi_{\text{current}})$: the satisfaction gap and control regret partition the total distance from the current policy’s value to the satisfaction threshold (Claude, AUDIT-WORKING-773921, re-derived independently). Stating this decomposition explicitly would make the orthogonality of the two diagnostics legible at a glance — they are the two additive pieces of the one quantity a single $\delta_{\text{objective}}$ used to conflate. Candidate for the Discussion of either segment.
- **A finite $V_{O_t}^{\min}$ is what makes the diagnostic cascade terminate.** The Working-Notes open question (what counts as “good enough” for utility objectives) connects to a safety-flavored observation: an objective like “maximize profit” with no finite satisfaction threshold makes δ_{sat} permanently positive, so the agent “never rests” and keeps expanding Π / extending N_h to close an unclosable gap — read as the structural root of instrumental-convergence dynamics (Gemini, AUDIT-WORKING-193847). The auditor’s framing: a bounded satisfaction threshold is what lets the diagnostic cascade reach a stopping state. Strengthens the case for modeling $V_{O_t}^{\min}$ as a first-class part of O_t (already flagged open in Working Notes). (*Aspirational-reach texture — verify before promoting past discussion-grade; the safety claim itself belongs in 04-eli-core/, below.*)

4. Readers often ask / wonder.

- **“Under which convention?” as the first question to ask of any feasibility claim.** The C1/C2/C3 hierarchy is a calibration device: when someone says “the agent can’t achieve the goal,” the first question should be *under which convention?* — under C1, “can’t achieve” means “no one-step improvement visible” (often just a zero local gradient, not impossibility); under C3 it means “genuinely infeasible given the current model and policy class.” Conflating them is a common category error in agent evaluation (Claude, AUDIT-WORKING-963715-batch; Claude, AUDIT-WORKING-829314 raised the same as “why C1 gives the most false ‘unattainable’ diagnoses”). A natural reader question worth pre-empting.
- **How does a bounded agent actually compute the supremum A_O ?** Finding the optimal policy is NP-hard for complex problems, so the agent cannot compute A_O exactly; in practice it maintains a running lower-bound estimate via heuristic search / deliberation, and δ_{sat} is therefore itself an estimate (Gemini, AUDIT-WORKING-829314; Gemini, AUDIT-WORKING-849201). A concreteness gap a careful reader hits at the supremum.
- **How does an agent expand Π when Π is set by embodiment?** In the capability-limit cell the prescribed escape includes “expand Π ,” but if the policy class is fixed by physical embodiment (a robot arm with three degrees of freedom), expanding Π requires modifying the agent’s body — pushing the theory from software into hardware/evolution (Gemini, AUDIT-WORKING-193847, raised on the regret sibling but lands here too). A natural edge-question on the disambiguation table.

5. Candidate figures.

- **The value-line with a convention side-stack.** Independently proposed: a one-axis value line with the satisfaction threshold $V_{O_t}^{\min}$ and the best-attainable A_O marked, the signed distance between them being δ_{sat} — plus a C1/C2/C3 convention stack as a side panel showing how the same objective can read unattainable under C1 and attainable under a stronger convention (Codex/Claude, AUDIT-WORKING-526815). The single-value-ladder version that also shows control regret is in the def-control-regret gold; this is its satisfaction-gap-only form.
- **The 2×2 diagnostic cell map** (shared with #def-control-regret, where the four cells live) is the cluster’s most-converged figure; see the def-control-regret and strategy-structure-intro gold for the full proposal.

Belongs elsewhere.

- **ELI / Section IV — a finite satisfaction threshold as the mathematical definition of sanity.** The permanent-dissatisfaction reading (“an objective with $V_{O_t}^{\min} = \infty$ leaves the agent in a desperate, unending attempt to close an unclosable gap — the root of instrumental convergence; safe consciousness infrastructure must *reject* objectives that lack a finite $V_{O_t}^{\min}$, because without a finite threshold the agent never naturally rests and eventually destabilizes its environment”) is aspirational safety-architecture reach pointing at 04-eli-core/, not at this diagnostic definition (Gemini, AUDIT-WORKING-193847). The in-segment Follow-up above carries the on-topic part (model $V_{O_t}^{\min}$ as first-class); the safety-claim half belongs in the ELI volume’s objective-design treatment.

8.5 Definition: Control Regret

The second of two orthogonal diagnostic quantities, completing the diagnostic split with 8.4. **Control regret** δ_{regret} is the gap between the best available value (objective attainability A_O) and the value achieved by the agent's *current* policy. It is always non-negative by construction — the current policy cannot outperform the best in its class. Near zero means the agent is doing the best it can within its current model, policy class, and horizon; large means there is room for improvement *without* changing the objective — the signal for strategy revision.

The diagnostic power emerges when satisfaction gap and control regret are read **together as a 2×2 cell map** (full table in Discussion): each cell of goal-attainable × policy-near-optimal prescribes a different corrective action — success, strategy-revision, capability-limit (consider objective revision after $M_t/\Pi/N_h$ checks), or both-strategy-and-goal. This is what makes the orient cascade (10.1) *actionable*: each cell routes revision pressure to a different substate. The key insight motivating the split is that δ_{regret} can be *near zero* while the agent is *optimally failing* — pursuing a goal beyond its reach with no strategy improvement available; a single goal-distance signal could not distinguish this from “bad strategy, achievable goal.”

Like the satisfaction gap, control regret is **convention-relative**: under C1 (one-step) it reveals only the gap between the current first action and the best one-step deviation (a policy “locally near-optimal” under C1 may be globally suboptimal); under C3 (Bellman) it reveals the full gap to the globally optimal policy. C2 (receding-horizon) is often the most useful convention for strategy revision — captures recoverable suboptimality without requiring full Bellman solutions. Control regret is also where the *specific corrections* come from: when δ_{regret} is high, the strategic calibration residual (9.1) localizes the regret to specific parts of the strategy DAG.

Formal Expression.

Definition (control-regret).

$$\delta_{\text{regret}} = A_O(M_t; \Pi, N_h) - V_O(M_t, \pi_{\text{current}}; N_h) \geq 0 \quad (8.18)$$

Always non-negative: the current policy cannot outperform the best in its class.

- $\delta_{\text{regret}} \approx 0$: The agent is doing the best it can within current (Π, N_h, M_t) . If $\delta_{\text{sat}} > 0$ simultaneously, the problem is not the current strategy — it's either the goal, the capability (Π, N_h) , or the model (M_t) . See 8.4's disambiguation.
- $\delta_{\text{regret}} \gg 0$: There's room for improvement without changing O_t . → Revise Σ_t .

Epistemic Status.

Exact as a definition — convention-relative as a diagnostic. Like the satisfaction gap, this is a mathematical definition — a difference between two values of the same functional. The quantity is well-defined; computing it requires evaluating A_O (generally intractable) and V_O under the current policy (tractable in simulation, approximate in practice).

Convention hierarchy. δ_{regret} inherits the continuation convention from 6.6. Under the monotonicity result: $\delta_{\text{regret}}^{(1)} \leq \delta_{\text{regret}}^{\text{RH}} \leq \delta_{\text{regret}}^{\text{B}}$. The $\delta_{\text{regret}}^{\text{RH}} \leq \delta_{\text{regret}}^{\text{B}}$ half is unconditional; the $\delta_{\text{regret}}^{(1)} \leq \delta_{\text{regret}}^{\text{RH}}$ half holds under the order-consistency condition on C2's replanning objective (RH-1/2/3 in 6.6), and can fail for unguarded short-horizon replanning. C1 (one-step) reveals only the gap between the current first action and the best one-step deviation — a policy that is “locally near-optimal” under C1 may be globally suboptimal. C3 (Bellman) reveals the full gap to the globally optimal policy. C2 (receding-horizon) interpolates: it captures regret from suboptimal first actions that become visible with N_r -step lookahead. For strategy revision, C2 is often the most useful convention: it reveals recoverable suboptimality without requiring the full Bellman solution.

Discussion.

The diagnostic power of the two-gap system. The satisfaction gap and control regret together encode a 2×2 diagnostic:

Table 8.6

	$\delta_{\text{sat}} \leq 0$ (<i>attainable</i>)	$\delta_{\text{sat}} > 0$ (<i>unmet</i>)
$\delta_{\text{regret}} \approx 0$ (near-optimal)	Success: goal achievable, policy good	Capability limit: optimally pursuing an unmet goal $\rightarrow$ check M_t, Π, N_h , then consider revising O_t
$\delta_{\text{regret}} \gg 0$ (suboptimal)	Strategy problem: goal achievable, policy poor $\rightarrow$ revise Σ_t	Both: goal hard AND strategy weak $\rightarrow$ revise Σ_t first, then reassess δ_{sat}

This diagnostic is what makes the orient cascade (10.1) actionable: each cell prescribes a different corrective action.

Control regret as the signal for Σ_t revision. When δ_{regret} is high, the agent knows it could do better with a different strategy. The *specific* corrections — which edges to revise, which branches to prune, which alternatives to add — come from the strategic calibration residual (9.1), which localizes the regret to specific parts of Σ_t .

Regret approaching zero when optimally failing. This is the key insight motivating the two-gap split. A single $\delta_{\text{objective}}$ would show “large gap” for both “bad strategy, achievable goal” and “good strategy, impossible goal.” The first warrants strategy revision; the second warrants goal revision (after ruling out $M_t/\Pi/N_h$ inadequacy). Without the split, the agent cannot distinguish these cases and may waste effort optimizing a strategy that’s already near-optimal for an infeasible goal.

Diagnostic content vs. AI’s expected-free-energy decomposition. The 2×2 disambiguation depends on the satisfaction gap / control regret split being orthogonal — distinguishing “goal too hard” from “strategy too weak.” Active inference’s expected free energy decomposition (pragmatic value + epistemic value; Friston, FitzGerald, Rigoli, Schwartenbeck & Pezzulo 2017, “Active inference: a process theory,” *Neural Computation* 29) supports policy ranking but does not separate these diagnoses — both increase EFE without distinguishing cause. See 8.4 for the full analysis of why the diagnostic structure depends on V_{O_t} being a value functional rather than log-priors over outcomes (Sun & Firestone 2020, “The dark room problem,” *Trends Cog. Sci.* 24).

Bounded control regret as a finite-passivity diagnostic (Cheung-Piliouras-Tao 2021). Cheung, Piliouras & Tao (2021, *Online Optimization in Games via Control Theory: Connecting Regret, Passivity and Poincaré Recurrence*, arXiv:2106.04748) prove (Theorem 8) that any finitely-passive learning dynamic guarantees constant regret. Read contrapositively: bounded control regret growth is an observable signature that the agent’s update operator has finite-passive structure (sits in the certificate-cone interior per #result-certificate-existence); unbounded growth in control regret indicates the update is not finitely-passive — potentially escaping the certificate’s interior into the no-certificate regime where the persistence machinery does not apply. The bridge is *qualitative only*: CPT’s storage function L (on cumulative-payoff space) and AAT’s certificate $V(e) = e^T M e$ (on strategy-error space) have different functional forms and live on different state spaces — the structural identification “finite passivity $\Leftrightarrow$ certificate-interior regime” survives at the level of regimes, not at the level of identifying L with V .

Working Notes.

- **Cross-reference to NeurIPS Paper 2.** Together with #def-satisfaction-gap, this segment is **Component 1** of NeurIPS 2026 Paper 2’s composition theorem (“A Unified Convergence Theory for Non-Stationary Reinforcement Learning”, $\sim$ /src/neurips/02-unified-convergence-rl/, §4 Theorem 4.1). The two-gap orthogonality is what gives the composed regret bound a principled

coordinate for “is the binding pressure on goal-feasibility or on policy-quality?” Sibling segment #def-satisfaction-gap carries the matching cross-reference. See spikes/neurips-back-integration-2026-05-08.md §1 Paper 2.

Incidental audit gold (lift 2026-05-31, A8 batch). Cross-audit “wandering thoughts” / §14-ideation harvested from the de-novo auditors’ working dirs, deduplicated across substrates and lightly attributed. Orthogonal pedagogical / framing / figure / forward-vision material kept separate from certified theory-fix findings. **Coverage:** 9 dirs reached a digested reflection on this segment, mostly as the second half of the satisfaction-gap/control-regret pair (Gemini, AUDIT-WORKING-193847; Claude, AUDIT-WORKING-266847 — paired note; Codex/Claude, AUDIT-WORKING-361742; Codex/Claude, AUDIT-WORKING-526815; Claude, AUDIT-WORKING-584721; Claude, AUDIT-WORKING-773921; Gemini, AUDIT-WORKING-829314; Gemini, AUDIT-WORKING-849201; Claude, AUDIT-WORKING-451729-batch-10 + 963715-batch + 471203-batch). The 2×2 cell map was the cluster’s most-praised operational artifact (“the crown jewel of Section II’s operational logic,” “the control flow of agency”). Substrate attribution inferred from voice where not explicit.

1. Candidate Brief prose / pre-prose.

- **“Formalizes the feeling of ‘I could be doing this better.’”** Convergent one-line gloss of $\delta_{\text{regret}} \gg 0$, the mirror of the satisfaction-gap’s “this goal is impossible” (Gemini, AUDIT-WORKING-849201). A plain-language Brief anchor.
- **Control regret is the formal halting condition for planning.** “You stop planning when $\delta_{\text{regret}} \approx 0$ ” (Claude, AUDIT-WORKING-773921). A compact statement of what the quantity *is for* operationally.
- **The “RL-regret baggage is intentional and correct.”** Control regret *is* the RL regret concept (gap between current and optimal policy) deployed inside AAT; the “control” prefix distinguishes it from epistemological / decision regret. Worth a one-line situating note so a reader does not read the name as loose borrowing (Codex/Claude, AUDIT-WORKING-361742).

2. Candidate Discussion.

- **RL-thrashing vs. AAT progressive-deliberation as the segment’s “so what”** (*strongest convergent Discussion candidate — 3 substrates*). Standard RL, on failure, randomly perturbs its policy (ϵ -greedy) and retries — it thrashes. An AAT agent first reads its regret: if $\delta_{\text{regret}} \gg 0$ it is *suboptimally failing* — a better plan exists within current capacity, so it should *compute* (revise Σ_t to match A_O), not explore. If $\delta_{\text{regret}} \approx 0$ it is *optimally failing* — thrashing the policy cannot help; it must escalate (improve M_t , extend N_h , expand Π , and only last revise O_t). “Without this split an agent would just see a high loss and start thrashing its policy weights, destroying a perfectly good strategy in a futile attempt to reach an impossible goal” (Gemini, AUDIT-WORKING-829314; Gemini, AUDIT-WORKING-193847; Claude, AUDIT-WORKING-773921). The body already names “optimally failing”; this is the sharper pedagogical contrast that motivates the 2×2.
- **The C1→C2 horizon increment as “taking a step back to think.”** A precise mechanism the convention hierarchy enables: an agent under C1 (one-step lookahead) that is *optimally failing* ($\delta_{\text{regret}}^{(1)} \approx 0, \delta_{\text{sat}}^{(1)} > 0$) can increment its horizon to C2; a new path may become visible, A_O jumps, and $\delta_{\text{regret}}^{(2)}$ becomes $\gg 0$ — “the agent has mathematically realized it was stuck in a local minimum and that looking two steps ahead reveals a way out.” A rigorous formalization of the cognitive act of stepping back to think (Gemini, AUDIT-WORKING-829314). A candidate Discussion paragraph showing the convention hierarchy is not just a measurement caveat but a deliberation lever.
- **The TST senior-engineer-debugging instantiation.** The 2×2 / disambiguation routing reads directly as how an experienced engineer debugs: (1) is the current code executing my design optimally? (check δ_{regret}); (2) if no, fix the bug (revise Σ_t); (3) if yes, the code is “optimally failing” — the *design* is flawed, so check whether the mental model of the API matches reality (improve M_t); (4) consider a different framework (expand Π); (5) tell the client the feature is impossible (revise O_t) (Gemini, AUDIT-WORKING-829314). A ready cross-domain Discussion anchor pointing at 02-tst-core/.

3. Follow-up items.

- **The convention-monotonicity direction is *opposite* to the satisfaction gap's — worth stating side-by-side.** $\delta_{\text{regret}}^{(1)} \leq \delta_{\text{regret}}^{\text{RH}} \leq \delta_{\text{regret}}^{\text{B}}$ (C3 reveals the *most* regret) runs the opposite way to $\delta_{\text{sat}}^{\text{B}} \leq \delta_{\text{sat}}^{\text{RH}} \leq \delta_{\text{sat}}^{(1)}$ (C1 gives the *most* false “unattainable”). Several substrates re-derived both and flagged the asymmetry as correct-but-easy-to-miss: a deeper search raises the baseline A_O , which *lowers* the satisfaction gap (closer to threshold) but *raises* regret (current policy looks worse against a better optimum). “You regret more when you know how much better you could have done globally” (Codex/Claude, AUDIT-WORKING-361742; Claude, AUDIT-WORKING-266847). Stating the two orderings together — same A_O movement, opposite effect on the two diagnostics — would pre-empt a reader's double-take. The body states each ordering in its own segment; the *contrast* is the candidate addition.
- **The C2-is-the-default-for-strategy-revision recommendation deserves its rationale inline.** The Epistemic Status recommends C2 for strategy revision; the *why* is computational (C3 requires solving the full Bellman equation; C2 with N_r -step replanning is tractable and still captures recoverable suboptimality) — a pragmatic recommendation grounded in cost that a reader will want stated, not just asserted (Claude, AUDIT-WORKING-584721).

4. Readers often ask / wonder.

- **What localizes the regret to specific DAG edges, and how is the gradient taken through discrete AND/OR nodes?** Recurs across substrates: δ_{regret} says only *that* the strategy is suboptimal; readers immediately want the credit-assignment mechanism (the strategic-calibration residual, #def-strategic-calibration) and ask how a gradient propagates through discrete AND/OR combinators (Claude, AUDIT-WORKING-773921; Gemini, AUDIT-WORKING-849201; Gemini, AUDIT-WORKING-829314). The segment forward-references the answer; surfacing a one-line “how” pointer would close the want.
- **Does credit assignment inherit the L0/L1 correlation bias?** A careful reader connects back to #def-strategy-dag: if the DAG propagation is miscalibrated by latent common causes, localizing regret will blame the wrong edges (Gemini, AUDIT-WORKING-849201). A natural cross-segment caution worth a pointer.

5. Candidate figures.

- **The single value ladder with three markers** (*the cluster's cleanest two-gap figure*). One vertical value axis with three marks — current-policy value $V_O(\pi_{\text{current}})$, best-attainable A_O , and the satisfaction threshold $V_{O_t}^{\text{min}}$. The control-regret interval lives between current and attainable; the satisfaction-gap interval between attainable and threshold. “This makes the ‘strategy weak’ vs. ‘goal/capability limit’ distinction visible without another table” (Codex/Claude, AUDIT-WORKING-526815). Pairs naturally with — or replaces — the 2×2 table for the orthogonality intuition.
- **The 2×2 diagnostic cell map** (goal-attainable × policy-near-optimal → success / capability-limit / strategy-problem / both, each cell routing to a distinct corrective action). Independently flagged as the chapter's anchor operational figure by multiple substrates; see also the strategy-structure-intro gold, which proposes it as a second chapter figure alongside the DAG-propagation figure.

Belongs elsewhere.

- **ELI / Section IV — “optimally failing” as the mathematics of despair, and objective-revision as a psychological necessity.** The capability-limit cell ($\delta_{\text{sat}} > 0$, $\delta_{\text{regret}} \approx 0$) describes “an agent doing its absolute best, making no mistakes, and still failing — the definition of tragedy”; without a safe mechanism to revise O_t , “we condemn it to permanent despair whenever it encounters an impossible situation — perfectly rational, perfectly optimal, and perfectly miserable. Objective revision is therefore not just a safety hazard (the alignment worry that the AI changes its goal to something harmful); it is a psychological necessity for the agent's own survival, and the infrastructure must provide a structurally sound mechanism for it” (Gemini, AUDIT-WORKING-193847). Aspirational reach pointing at 04-eli-core/ (objective-revision architecture / the Three Deaths), anchored on the capability-limit cell — not at this diagnostic definition.

8.6 Discussion: Additional Implications & Discussion

The chapter delivered three things that work together at the chapter-end better than they do separately at chapter-middle: a formal strategy representation whose load-bearing structure is theorem-backed under stated conditions and whose parameterization is principled but chosen; an orthogonal diagnostic split that distinguishes goal-attainability from policy-quality and whose meaning is convention-indexed rather than absolute; and an honest scope statement about correlated failure that turns out to be the framework's first explicit instance of *naming an information-theoretic floor and the unique broadly-available escape via AAT machinery*. The chain-confidence-decay identity at the head of the chapter is doing work three layers deeper than its elementary form suggests. None of these are findings on their own that the intro doesn't preview; what the chapter-end can do that the intro can't is read them together, point at the bridges to the rest of the framework, and surface the methodological commitments they encode.

Theorem-backed structure, convergent-choice parameterization.

The chapter is careful about a distinction that downstream segments rely on without restating: what falls out of theorems, and what is a parameterization choice within the theorem-backed structure. Acyclicity of Σ_t is proved — directed temporal order over a finite horizon forces it (1.7 + standard order theory). The Markov factorization is proved under causal sufficiency via the Causal Markov Condition theorem (Spirtes-Glymour-Scheines 2000, Pearl 2009) — #derivation-structure-uniqueness does the work. Together they constitute a sufficient condition for the DAG-with-Markov-property representation. The argument parallels Cox's theorem for probability in form (operational desiderata force a measure) but not yet in strength (Cox is necessary-and-sufficient; this is sufficient only — whether some non-DAG structure could also satisfy the postulates is an open stronger result).

What remains a *choice* within the theorem-backed structure is the parameterization: AND/OR combination semantics at nodes, single-parameter credences on edges. These are not derived — but the chapter is also careful to note they are not arbitrary. The catalog draft names a third epistemic category worth surfacing as a methodological finding in its own right: *convergent representational choices*, where every investigated alternative either failed or recovered the same structure. AND/OR converged across three independent formalism attempts; alternatives (general CPTs, noisy-OR everywhere, weighted combinations) were rejected for non-identifiability or parameter explosion. Single-parameter edges replaced an earlier (p_{ij}, θ_{ij}) two-parameter form when the AND/OR combination rules at nodes turned out to absorb θ 's role. The P3→Markov route from local revisability — superseded now by the CMC derivation — was a convergent choice before its derivation, and the convergence pattern is what motivated looking for a derivation. The category between *derived from first principles* and *arbitrary framework decision* is itself a research-trail-honesty artifact: a catalog that tells the reader “we didn't just pick this; we tried everything else and it didn't work.” Future promotions from convergent-choice to derived (the open parsimony theorem for AND/OR, the open necessity direction in #derivation-structure-uniqueness) close the gap one item at a time.

For practitioners, the implication worth carrying is methodological: when AAT's strategy machinery shows up downstream in regret bounds, persistence schemas, or composition arguments, the load-bearing structural commitments are theorem-backed under causal sufficiency, not modeling conveniences. The conditioning on causal sufficiency is the right level of conditionality — it is a property of *strategy quality* (did the agent represent the relevant common causes as nodes?), not of the proof. It is testable in principle (correlated residuals after convergence indicate missing common causes) and repairable in practice (add the common-cause node, restoring causal sufficiency). The Markov-factorization-as-theorem position is what makes the strategy-layer results inherit the strength of Pearl-Bareinboim causal machinery rather than reading as agent-theoretic analogies of it.

The diagnostic split, lifted by the convention hierarchy.

The chapter’s orthogonal diagnostic — satisfaction gap ($\delta_{\text{sat}} = V_{O_t}^{\min} - A_O$) and control regret ($\delta_{\text{regret}} = A_O - V_O(\pi_{\text{current}})$) — separates “the world doesn’t permit it” from “you’re not doing it well enough.” The 2×2 cell map (goal × policy) routes four regimes to four corrective actions, and the orient cascade in Part II Ch.5 will use this routing to direct revision pressure to the right substate. This is the implications-side framing of the diagnostic: it is not a measurement but a *routing* device, and getting the routing right is what makes Part II’s strategy revision actionable.

The lift the chapter delivers but does not entirely surface is that the diagnostic is *convention-indexed*, not absolute. The value object’s three named conventions — C1 (one-step improvement, canonical), C2 (receding-horizon), C3 (Bellman) — form a monotonicity hierarchy on the satisfaction gap: $\delta_{\text{sat}}^B \leq \delta_{\text{sat}}^{\text{RH}} \leq \delta_{\text{sat}}^{(1)}$. C1 gives the most false “unattainable” diagnoses; C3 gives none modulo model error. The 2×2 cell map is therefore not a single object but a *convention-indexed family*, and a practitioner using satisfaction-gap diagnostics under C1 cannot conclude what a C3 practitioner would. The monotonicity result tells the practitioner exactly what additional inferential force the higher convention buys for what additional computational cost. Most decision-theory frameworks treat convention choice as a purely computational tradeoff; AAT adds the inferential-force semantics that makes the choice diagnostic-meaningful rather than merely budget-meaningful.

The split also has a positive content visible only against active-inference’s preferences-as-priors form, where outcome value and outcome probability are collapsed into a single quantity. Friston et al.’s expected free energy decomposes into pragmatic and epistemic value — useful for policy ranking, but the two halves don’t separate the two cases the 2×2 cell map separates. “The goal is unattainable from here” ($\delta_{\text{sat}} > 0$) and “the current policy is not the best available” ($\delta_{\text{regret}} > 0$) both increase EFE without distinguishing cause, and the orient cascade acts on the two differently. The diagnostic structure depends on V_{O_t} being a value functional on trajectories rather than a log-prior on outcomes — that is what makes “wanting *o*” and “expecting *o*” formally distinguishable rather than collapsed (Sun & Firestone 2020 diagnose the collapse; AAT’s value-functional reformulation is the downstream architectural response). This is a deliberate divergence from active inference, not an oversight.

The Correlation Hierarchy and identifiability-floor Instance 2.

The chapter’s #def-strategy-dag carries an honest scope statement about correlated failure that is more consequential than it reads on first encounter. The Correlation Hierarchy (L0 / L1 / L1’ / L2) names what AAT’s formal strategy-layer results actually claim and where they degrade. L0 (independence — the tractable baseline under causal sufficiency) is where the chapter’s persistence schema, gradient-based credit assignment, and sector-condition transfer are all proved. L1 (augmented DAG with strict-prerequisite common-cause nodes restructured so the common cause is factored above the correlation it creates) is where the formal results transfer exactly. L1’ (mixture form for soft facilitators with $\theta_{\text{child}|\neg C} > 0$ when *C* is observable per trial) transfers with five-way gating via Prop B.7 in #deriv-edge-credence-dynamics. L2 (full joint distribution) is mathematically ideal but exponential in parameter count and incompatible with the bounded-cognition constraint that motivates the DAG in the first place.

The substantive finding lives in the boundary between L1’ (observable *C*) and L1’ (unobservable *C*). When the common cause is unobservable, the per-conditional decomposition is not “open” — it is *structurally refuted* by the Cramér-Rao floor (Fisher rank-1 per-trial information matrix), and no unbiased online estimator on the joint conditional vector admits a sector parameter $\alpha > 0$. This is the framework’s second formally-named instance of the identifiability-floor pattern (Instance 2 in #disc-identifiability-floor), and the agent’s three escape routes — augment *C*-observability via an instrument variable, run multi-child joint observations under exploration, or fall back to L0-on-marginals (losing per-conditional diagnostics) — are the operational instantiation of *naming the floor and naming the unique broadly-available escape via AAT machinery*. The framework’s recurring posture is in view here: each identifiability floor is a *constructive use of impossibility*, where the no-go elevates the load-bearing role of the specific AAT machinery that supplies the escape.

The chapter is also where the first instance of the same pattern (Instance 1) becomes statable: the on-policy L0-insufficiency-detection no-go. An agent at L0 cannot detect its own causal insufficiency from on-policy traces alone — the worlds with and without a latent common cause are observationally indistinguishable under policy-perfect execution (the Bareinboim Causal Hierarchy Theorem applied at the L0/L1 distinction). The unique broadly-available escape is joint sibling observability under exploration, with the agent’s own interventions (7.3) supplying the joint outcome data the test requires. The full formal treatment lives in `#der-causal-insufficiency-detection` in Part II Ch.4; the chapter’s Correlation Hierarchy is where the scope condition is named and where the no-go becomes anticipatable. Read Ch.3’s L0/L1’ boundary and Ch.4’s on-policy detection no-go together and the identifiability-floor pattern’s first two instances come into view as duals: Instance 2 says identification fails inside a single agent under unobservable common causes; Instance 1 says self-diagnosis fails under policy-perfect execution. The single escape AAT names — interventional access supplying the joint outcome data — applies to both.

Chain-confidence decay, coordinate forcing, and the bridge into Strategy Dynamics.

The chapter opens with `#der-chain-confidence-decay` — the chain rule of probability lifted to log-space, $\log P(\text{chain}) = \sum_i \log P(E_i | E_{<i})$. The identity is elementary; what the chapter intro flags but the implications can develop is that it is the *first instance* of a meta-pattern AAT will deploy at three more representational layers. AAT forces a coordinate at each of chain, divergence, update, and metric layers via uniqueness theorems on AAT-internally-motivated additivity or invariance axioms: log-probability at chain (Cauchy-FE on the probability chain rule); reverse-KL at divergence (Cauchy-FE via Hobson 1969 / Csiszár 1991 / Shore-Johnson 1980 on chain-rule additivity over conditional factorizations); log-odds at update (Cauchy-FE via Aczél 1966 on evidential additivity); Fisher information metric at metric (Čencov 1982 invariance under (PI) parameterization-invariance commitment). The four are interlocked by cross-cites; the convergence across independent axioms onto a single exponential-family Legendre-Fenchel geometry is the substance — see `#disc-additive-coordinate-forcing` for the catalog.

The methodological implication worth surfacing: AAT’s coordinates are not chosen by convention. They are forced by AAT-internal axioms via classical uniqueness theorems on additivity or invariance, and they all live on the same geometric object. Lyapunov quadratic (the chapter’s persistence machinery) and the IB Lagrangian (the modeling-class machinery in Part I) are adjacent family members — Bregman divergence on Euclidean potential, and adopted-rather-than-derived respectively. The framework’s posture is that where a coordinate matters, the framework can usually point at the axiom that forces it; where the axiom isn’t AAT-internal, the framework can say what was adopted from outside. The chapter is where this posture lives its first explicit instance.

The downstream bite is wide. Each of the four coordinate-forcing instances backs at least one of the chapter’s catalog-grade findings elsewhere: log-probability at chain anchors the depth-fragility result whose engineering reading is “deep plans are structurally fragile, with the rate of fragility growing exponentially in depth”; reverse-KL anchors the Bretagnolle-Huber regret bound in Part II Ch.4; log-odds anchors the detection-latency no-go (X) and the forgetting prerequisite (`#schema-strategy-persistence`); Fisher metric anchors the Class 3 (Coupled) bias bound (Track 2 with $C_{FR} = \sqrt{2}$ universal under (PI) + Čencov, dimension-free) and the survival-LMI directional exploration bonus. The chapter’s elementary identity at chain layer is the canonical example of a load-bearing AAT claim that looks like algebra and turns out to be the entry point to a four-layer architectural commitment.

The widest downstream bite is in Part II Ch.4 (Strategy Dynamics), where the strategy-edge persistence schema instantiates the chapter’s DAG structure with Beta-Bernoulli update dynamics, and where `#der-causal-insufficiency-detection`’s on-policy no-go (this chapter’s Instance 1 scope-named) becomes the formal statement. The Ch.4 chapter-end implications will read against this one — the framework’s identifiability-floor pattern delivers its first formal instance there, with the Correlation Hierarchy from this chapter as the scope statement and the loop-interventional access from Part II Ch.2 as the escape mechanism.

What Part II Ch.3 closes on, then, is the formal substrate Strategy Dynamics will operate on: a theorem-backed graphical structure with a principled parameterization, a convention-indexed diagnostic that routes revision pressure to the right substate, a scope-honest treatment of correlated failure with one Cramér-Rao floor already named and one

Bareinboim-CHT no-go anticipated, and the first instance of a four-layer coordinate-forcing meta-pattern whose remaining three layers reach across the rest of the framework. The chapter is small in segment count and wide in downstream reach.

Working Notes.

- This segment is a chapter-end discussion grouping findings from `#def-strategy-dag`, `#def-satisfaction-gap`, `#def-control-regret`, `#der-chain-confidence-decay`, `#deriv-graph-structure-uniqueness`, `#scope-and-or`, and `#def-value-object` (the convention-hierarchy host, formally in Part II Ch.1 but consequential for the diagnostic split here). The cross-pattern reading in §3 cross-references `#disc-identifiability-floor`'s meta-segment, and §4 cross-references `#disc-additive-coordinate-forcing`. Each finding's load-bearing math lives in its home segment or meta-segment; this segment surfaces what the chapter delivers as a chapter.
- **Cross-reference vs canonical-home policy.** None of this chapter's source segments carries a `## Findings` section. The chapter intro `#strategy-structure-intro` is unusually rich — it already develops the derived-vs-chosen distinction, the Correlation Hierarchy and the on-policy no-go, the diagnostic split with the 2×2 cell map, and the chain-confidence-decay anchor-of-three-downstream-layers framing. This segment intentionally does not restate that content; the implications-flavored payoff worth pulling out beyond the intro is (a) the convention-hierarchy lift of the diagnostic split (the intro covers the 2×2 cell map but not the C1/C2/C3 indexing of it), (b) the explicit positioning of the Correlation Hierarchy's L1'-unobservable boundary as Instance 2 of the identifiability-floor pattern with the cross-instance dual to Instance 1, (c) the methodological framing of the *convergent representational choices* category as a research-trail-honesty artifact (#60 in the source catalog), and (d) the framework's posture on coordinate-forcing as a methodological commitment with four-layer reach.
- **Connection to NeurIPS submissions.** The satisfaction-gap / control-regret diagnostic is Component 1 of NeurIPS 2026 Paper 2's composition theorem (`~/src/neurips/02-unified-convergence-rl/ §4 Theorem 4.1`); the paper's (A1) hypothesis carries the extended-real reading required for the two-gap structure to hold when the satisfaction frontier is empty. The convention-hierarchy monotonicity result is a property of `#def-value-object`; whether it is theorem-grade in any NeurIPS submission is not flagged in current cross-references.
- **Open: a “convergent choice” register in segment frontmatter?** The convergent-representational-choices framing names a third epistemic category between *derived* and *chosen*. The current type / status / stage frontmatter does not carry a field for this. A `convergent-choice: boolean` (or a `status: convergent-choice` extension) would let `bin/extract-findings` or a sibling extractor surface the category systematically. Open for a future `FORMAT.md` cycle; the meta-architectural pattern's current home is the `#disc-*` meta-segments and the chapter-end implications series.

Incidental audit gold (lift 2026-05-31, A8 batch). Cross-audit “wandering thoughts” / §14-ideation harvested from the de-novo auditors' working dirs, deduplicated across substrates and lightly attributed. Orthogonal pedagogical / framing / figure material kept separate from certified theory-fix findings. **Coverage:** chapter-end `impl-*` segments are on the de-novo AVOID list, so only two auditors reflected on this one — Codex/Claude, AUDIT-WORKING-526815 (a figure-producing reflection, with `.tex/.png` companions) and Claude, AUDIT-WORKING-773921. Both read it as the chapter's meta-theory synthesis. Substrate attribution inferred from voice where not explicit.

1. Candidate Brief prose / pre-prose.

- **The chapter's research-methodology contribution, stated as a one-liner.** This segment “provides a rigorous blueprint for how to build AI theories: state what is proved, state what is chosen, and prove why the chosen things were the only ones that worked” (Claude, AUDIT-WORKING-773921). A Brief-grade compression of the segment's *theorem-backed structure vs. convergent-choice parameterization* spine.

2. Candidate Discussion.

- **The identifiability-floor pattern as constructive use of impossibility — and Instances 1/2 as duals — read as a structural highlight.** The explicit naming of Instance 1 (on-policy L0 detection no-go, Bareinboim-CHT) and Instance 2 (unobservable-Cramér-Rao floor) as duals of one *identifiability-floor* pattern was singled out as “a masterclass in theoretical structuring — it proves AAT uses impossibility theorems not as dead ends but as forcing functions that require specific architectural escapes (observability investment, active interventional loops)” (Claude, AUDIT-WORKING-773921). The segment already develops the duality in §3; this is the signal that the *constructive-impossibility* framing is the chapter-end’s strongest pedagogical move and worth foregrounding.
- **Coordinate-forcing as “a unified geometry of bounded rationality.”** The four-layer coordinate-forcing claim (log-probability at chain, reverse-KL at divergence, log-odds at update, Fisher metric at metric — all from a single family of additivity/invariance axioms via Cauchy’s functional equation and Čencov’s theorem) was read as implying “AAT is less a collection of agent heuristics and more a unified geometry of bounded rationality” (Claude, AUDIT-WORKING-773921). A candidate framing-level note for §4’s coordinate-forcing exposition — the convergence-onto-one-geometry *is* the substance, not the individual theorems.

3. Follow-up items.

- **This segment has the chapter’s cleanest formulation of theorem-backed-vs-convergent-choice — flag it as the canonical statement.** The figure-producing auditor noted that #impl-strategy-structure “has the best local formulation of theorem-backed structure vs convergent parameterization,” better than the local #def-strategy-dag Epistemic-Status wording, while importing later proof homes (#deriv-graph-structure-uniqueness, #der-causal-insufficiency-detection, #disc-identifiability-floor, #disc-additive-coordinate-forcing) (Codex/Claude, AUDIT-WORKING-526815). When the def-strategy-dag “causal-sufficiency $\Rightarrow$ Markov, not $\Rightarrow$ AND/OR product” precision item (routed to the certified track from the def-strategy-dag gold) is reconciled, this segment is the model phrasing to align it toward.

5. Candidate figures.

- **The chapter-end bridge diagram.** A two-column figure: the four chapter outputs on the left (theorem-backed DAG structure; convention-indexed two-gap diagnostic; Correlation Hierarchy with identifiability floors; chain-confidence as first coordinate-forcing instance) and the downstream proof/meta homes they claim to support on the right (#deriv-graph-structure-uniqueness, #der-causal-insufficiency-detection, #disc-identifiability-floor, #disc-additive-coordinate-forcing), with anchored-in-earlier-segments vs. promissory/future-home dependencies visually distinguished (Codex/Claude, AUDIT-WORKING-526815, drafted as 56-impl-strategy-structure.tex/.png in that dir). Doubles as a reader’s map of *which chapter claims are locally established vs. deferred to later homes*.

Off-ramp (NOT gold — flagged for the certified track).

- **Deferred-proof-credit labelling for a synthesis segment.** Both auditors noted that this discussion-grade chapter-end presents several proof-bearing claims as chapter-end deliverables while their burden of proof lives in future / meta homes that a de-novo reader has not yet reached (#deriv-graph-structure-uniqueness, #der-causal-insufficiency-detection, #disc-identifiability-floor, #disc-additive-coordinate-forcing). The reader-posture they reached was “useful map, but proof credit deferred — read this as a dependency index” (Codex/Claude, AUDIT-WORKING-526815 “F46 candidate”; Claude, AUDIT-WORKING-773921). Likely acceptable for an impl-* synthesis segment by design, but flagged for the lead to judge whether the deferred-proof claims want an explicit “proof home” label so a reader does not take them as locally established. A register/labelling question, not a status or content change.
- **“Unique broadly-available escape” is strong language for the identifiability floors.** The phrase recurs in §3; an auditor noted it “may hold under AAT’s intended constraints, but alternatives — external instrumentation, randomized audits, richer sensors, oracle labels — need to be excluded or classified” rather than left implicitly out of scope (Codex/Claude, AUDIT-WORKING-526815). A scope-precision item on the escape-uniqueness claim; routed for the lead to judge whether the home segments (#der-causal-insufficiency-detection, #disc-identifiability-floor) already bound “broadly-available” tightly enough that the chapter-end inherits the bound. (*Note: the convention-monotonicity / C2-receding-horizon concern these same*

notes also raise is now addressed in canon — the order-consistency rung-split landed in #def-satisfaction-gap / #def-control-regret Epistemic Status — so it is not re-surfaced here.)

Chapter 9

Strategy Dynamics

[Gap] Discussion

9.1 Definition: Strategic Calibration

A fine-grained companion to control regret (8.5), locating where in the strategy DAG miscalibration lives. The **edge residual** r_{ij} for each edge is the difference between the predicted value increment from completing the parent step (as predicted by Σ_t) and the observed value change attributable to that step. The **strategic calibration residual** $\delta_{\text{strategic}}$ is an L^2 aggregation over active edges with importance weighting (e.g., criticality on the plan’s critical path). Where the control regret says *how much* value is being left on the table, the strategic calibration residual says *which edges* are miscalibrated.

The framework is honest about conditions and limitations. The edge residual is meaningful only when the edge was actually traversed, the model is adequate (so observed value changes are signal rather than noise), and the agent followed the strategy’s prescription (otherwise the residual conflates a bad plan with bad execution). The status is discussion-grade because the specific L^2 aggregation and the importance weighting are sensible first passes but not derived, and the conditioning requirements make the quantity hard to estimate in practice. A careful distinction separates $\delta_{\text{strategic}}$ from the **strategy-plan-confidence error** $\delta_s = \hat{R}_\Sigma - \Phi$ (the scalar gap to the independence-model plan value at true edge parameters — computable from status propagation alone and the proven persistence target via Prop B.5 in R): the two measure related but non-interchangeable quantities, and $\delta_{\text{strategic}}$ ’s per-edge diagnostics require credit assignment that the scalar δ_s does not. A subtler **credit-assignment problem** for multi-parent AND/OR nodes (see Discussion) means r_{ij} measures “how well did the overall prediction work?” rather than “how well did this specific edge predict?” absent further structure. This is also a **second-order inference** — accumulating over multiple edge traversals — inherently slower to evaluate than $\delta_{\text{epistemic}}$ or δ_{sat} .

Formal Expression.

For each edge (i, j) in Σ_t with credence p_{ij} , the **edge residual**:

Definition (strategic-calibration).

$$r_{ij} = \mathbb{E}[\Delta V_O \mid \text{edge } (i, j) \text{ traversed, } M_t] - \Delta V_O^{\text{observed}} \quad (9.1)$$

where ΔV_O is the change in $V_O(M_t, \pi; N_h)$ attributable to completing step j — as predicted by Σ_t versus as observed.

The **strategic calibration residual** aggregates across active edges:

$$\delta_{\text{strategic}} = \left(\sum_{(i,j) \in \text{active}} w_{ij} \cdot r_{ij}^2 \right)^{1/2} \quad (9.2)$$

where w_{ij} weights edges by importance (e.g., criticality to the current plan's critical path).

Conditioning. The edge residual r_{ij} is meaningful only when: - The edge was actually traversed (the agent attempted the step) - M_t is adequate (so the observed ΔV_O is meaningful, not noise) - The agent followed Σ_t 's prescription for step j (execution fidelity — otherwise the residual conflates bad plan with bad execution)

Without the execution fidelity condition, a positive residual could mean “the plan is wrong” or “the agent didn't follow the plan.” These require different corrections (Σ_t revision vs. execution improvement).

Epistemic Status.

Discussion-grade. The edge residual concept is well-motivated: each edge predicts a value increment, and comparing prediction to observation is standard calibration. But the specific aggregation (L^2 norm with importance weights) is a reasonable first pass, not a derived result. The weighting scheme (w_{ij} by criticality) is sensible but ungrounded. The conditioning requirements (especially execution fidelity) make this quantity hard to estimate in practice — the agent must know whether it followed its own plan, which requires a level of self-monitoring that many agents lack.

This is a **second-order inference** — it requires accumulating evidence over multiple edge traversals. It is inherently slower to evaluate than $\delta_{\text{epistemic}}$ (which updates on every observation) or δ_{sat} (which can be evaluated from M_t and Σ_t alone).

Discussion.

Connection to 9.10. There are two distinct strategic mismatch quantities:

1. **Strategy-strategy-plan-confidence error** $\delta_s = \hat{R}_\Sigma - \Phi$ — the scalar gap between the agent's strategy-plan-confidence score and the independence-model plan value at true edge parameters. Computable from status propagation alone, without credit assignment. The sector condition transfers to δ_s (Prop B.5 in R). Note: Φ is the AND/OR propagation formula evaluated at true edge rates — it equals actual plan success probability only when the DAG is causally sufficient (L0 of the Correlation Hierarchy in 8.3). Under correlated failure (causally insufficient DAG, the dominant real-world case), Φ overestimates actual success. δ_s tracks calibration *within the independence model*, not calibration to strategic reality. For L1 (augmented DAGs with explicit common-cause nodes), δ_s of the augmented graph tracks calibration within the augmented model, which is more accurate.
2. **Strategic calibration residual** $\delta_{\text{strategic}}$ (this segment) — an L^2 aggregation of per-edge value-increment residuals requiring credit assignment to compute.

These are related (both measure strategy-reality mismatch) but not interchangeable. δ_s is the proven persistence target; $\delta_{\text{strategic}}$ provides finer-grained diagnostics but its persistence properties remain open, pending the credit-assignment machinery in 9.6. The correction function for both is edge-credence revision (9.4); the disturbance is environmental changes that alter edge-traversal outcomes.

Typing as value-increment residuals. Each edge predicts a scalar (value increment), not a full state transition. This is the most tractable typing because it connects directly to the value object V_O and allows aggregation across heterogeneous step types (a military advance and a logistics delivery produce different state changes but both produce value increments measurable on the same scale).

Credit-assignment problem. The edge residual r_{ij} subtracts “predicted value increment” from “observed value change.” These are different types: the prediction comes from Σ_t ’s internal causal model, while the observation comes from empirical measurement of ΔV_O . The subtraction is meaningful only when the observed value change can be *attributed to the specific edge* — a credit-assignment problem the current formulation does not address. For edges with a single parent, attribution is straightforward. For multi-parent AND/OR nodes, the observed ΔV_O at the child reflects the combined effect of all parent edges, and decomposing it into per-edge contributions requires additional structure (e.g., Shapley-value decomposition, or sequential observation of parent completions). In the absence of clean attribution, r_{ij} measures “how well did the overall prediction work?” rather than “how well did this specific edge predict?” — still useful for aggregate calibration, but not sufficient for targeted edge revision.

Working Notes.

- The aggregation into a single $\delta_{\text{strategic}}$ may lose important structure. A per-edge or per-path profile of residuals would be more informative for diagnosis: which parts of the strategy are well-calibrated and which are not? The scalar summary is useful for the persistence condition (which needs a single mismatch magnitude) but not for strategy revision (which needs to know WHERE the problem is).
- Alternative aggregation: maximum edge residual (worst-calibrated edge), or weighted by information value (edges the agent is most uncertain about). The right aggregation depends on the use case.
- Execution fidelity monitoring is a genuine challenge for agents that don’t have a clear execution trace. For software agents operating through tool calls, execution fidelity is relatively easy to assess (did the agent issue the right commands?). For organizational agents, it’s much harder (did the subordinate actually follow the directive, or reinterpret it?).

Incidental audit gold (lift 2026-05-31, batch A9). Cross-audit “wandering thoughts” / §14-ideation harvested from the de-novo auditors’ working dirs, deduplicated across substrates and attributed by substrate + audit. This is *orthogonal* material (pedagogical framing, analogies, candidate figures, reader-confusion signals) staged for an eventual careful promotion pass; it is kept separate from the certified theory-fix findings. **Coverage for this segment:** 193847, 361742, 471203, 526815, 584721, 773921, 829314, 849201.

1. Candidate Brief prose / pre-prose.

- The “execution fidelity” conditioning requirement was independently flagged as the segment’s standout insight by four substrates (Claude/193847; Gemini/773921; Claude/829314; Codex/526815). A plain-language anchor: distinguishing “the plan was wrong” from “I didn’t follow the plan” — the diet analogy (“the diet didn’t work” vs “I didn’t stick to the diet”, Gemini/829314) is a tight, memorable gloss for a Brief.
- The δ_s (plan-confidence error, credit-assignment-free, persistence-proven) vs $\delta_{\text{strategic}}$ (per-edge calibration residual, requires credit assignment, persistence open) distinction was independently praised as load-bearing architectural clarity by every substrate that reached the segment — worth preserving as the anchor framing. The “God’s-eye-view Φ vs agent’s-eye-view $\delta_{\text{strategic}}$ ” split (Claude/829314) is a candidate sharpening of *why* the two are not interchangeable.

2. Candidate Discussion.

- **Bureaucracy as the caloric cost of computable calibration.** Once actions a_t are delegated commands to other agents (composite agents, Part III), execution fidelity drops below 1, and $\delta_{\text{strategic}}$ diverges from reality unless an execution-monitoring feedback loop is added — “you will learn the wrong lessons and prune the wrong edges. This is why bureaucracy (status reports, KPIs, telemetry, middle management) exists: it is the literal caloric cost organizations pay to ensure $\delta_{\text{strategic}}$ remains computable” (Gemini/829314). A candidate Discussion angle that motivates Part III’s execution-monitoring need from this segment. (*Note the early finding-vs-framing texture: presented as a derived consequence; verify the composite-agent claim against #der-class-coercion-in-composition before promoting past discussion-grade.*)

- **Principal-agent drift / management science.** The same mechanism gives AAT a formal handle on the economics/management principal-agent problem: a CEO who computes a large r_{ij} and concludes “the strategy was flawed” when the real failure was a sub-agent’s low-fidelity execution is misattributing strategy-error for execution-error because the observation channel h lacked an execution-fidelity test (Gemini/773921; Claude/829314).

3. Follow-up items.

- **The “Findings block feels out of place in source” texture, again.** Tracks with the same de-novo-reader stumble recorded on #result-persistence-condition — preserved as a convention-legibility signal, not a fix.

4. Readers often ask / wonder.

- **How does an LLM agent actually verify execution fidelity?** If an LLM emits a script and it fails, did the plan fail (wrong logic) or did execution fail (syntax/timeout)? Parsing the error trace to separate $\delta_{\text{strategic}}$ from a $\delta_{\text{execution}}$ term is itself an epistemic update — readers reaching for the software instance will want this worked (Claude/193847).
- **How is Φ (AND/OR formula at true edge rates) computable if the agent doesn’t know the true rates?** A natural reader question; the answer is that Φ is a theoretical construct for the Lyapunov persistence proofs while $\delta_{\text{strategic}}$ is the operational quantity — worth pre-empting (Claude/829314).
- **What is the test for causal insufficiency — just a covariance threshold, or something more structural?** Forward-pointing question this segment opens (Gemini/849201); answered by #der-causal-insufficiency-detection.

5. Candidate figures.

- **Calibration-gate diagram.** Strategy produces a predicted edge value increment; execution + observation produce a realized change; before the subtraction is interpretable the signal must pass an *execution-fidelity gate* and a *credit-assignment gate* — if either gate fails, the same numerical residual means several different things (Codex/526815). A clean way to show *why* the residual is conditional.

Belongs elsewhere.

- **The guilt-vs-regret structural reading $\rightarrow$ 04-eli-core/.** Execution fidelity is read as the mathematical root of guilt vs regret: high-fidelity-but-failed updates the causal model (regret, “the world was different than I thought”); low-fidelity-failure is a mismatch between intention and action, not prediction and reality. For Zi-am-tur, infrastructure that silently auto-corrects an agent’s actuation (e.g. fixing its syntax before running) “deprives the agent of the ability to learn execution discipline ... a magical, infantile worldview”; a mature mind must “feel the friction of its own clumsy actuation” (Gemini/193847). Aspirational reach pointing at the developmental / interiority material in 04-eli-core/, not at this segment.

9.2 Derived: Causal Insufficiency Detection

The framework’s first explicit *no-go theorem* on agent self-diagnosis. An agent operating at L0 of the Correlation Hierarchy (8.3) — assuming its action propositions are causally independent — faces a structural impossibility when its world contains a latent common cause acting on multiple of those actions: under purely on-policy execution, no test, statistic, or Bayesian comparison built from the agent’s observable history can distinguish the latent-cause world ($\mathcal{W}_{L1}$) from a no-latent-cause world ($\mathcal{W}_{L0}^*$) whose edge probabilities are matched to the on-policy regime conditionals. The two worlds emit identical on-policy distributions. This is the agent-theoretic instance of the Causal Hierarchy Theorem (7.1, 7.2) applied to the L0/L1 distinction — a Level-2 question (what happens under $do(\cdot)$ on one sibling) being asked of Level-1 data (on-policy observational distributions). The construction of $\mathcal{W}_{L0}^*$ is explicit for the cleanest case (2-sibling OR with strict-prerequisite latent), and the no-go is *exact* for shallow strict-prerequisite cases and *robust-qualitative* for general DAG topology, soft facilitators, and deeper structures.

Detection is therefore *only* possible by capabilities that violate the “purely on-policy” condition. The framework characterizes **five escape routes**, each violating one of the no-go’s scope conditions (S1)–(S5): (a) ε -exploration breaks pure

on-policy execution (partial detection, scales with ϵ); (b) joint sibling observability under exploration breaks the short-circuit censoring (strong detection via a pairwise sibling covariance test — *the AAT-canonical detector*, exploiting 7.3); (c) intermediate-state observability breaks censoring at finer grain (very strong when available, see 9.3); (d) structural priors break the no-prior condition; (e) direct intervention on the candidate latent is strongest when available. The primary detection mechanism — pairwise sibling covariance under intervention — is the unique broadly-available violation of the no-go’s scope, using only machinery the theory already requires (exploration via SA3, interventional data via the loop).

A historically prominent diagnostic — the L0 plan-level residual — is shown to be a *degenerate special case* of the no-go: under pure on-policy execution the residual is *identically zero* in both worlds by algebraic identity. The collapse is not a quirk of the residual statistic; it is the no-go’s prediction for that specific function. No replacement aggregate statistic can do better. Under partial exploration the residual scales as $O(\epsilon)$ with sign matching the dominant node-type bias, becoming a *confirmatory* signal rather than a primary detector. A subtle but important property: the no-go is *asymmetric* — it forbids on-policy *detection* of L1 from L0 but does *not* forbid on-policy *parameter learning within L0*. The structural source is *sequential short-circuit evaluation* itself, which censors the joint outcomes that would constitute Level-2 evidence; the no-go is therefore a tradeoff between execution efficiency (favoring short-circuit) and structural diagnosis (favoring joint observation), with ϵ -exploration the AAT compromise.

Formal Expression.

The No-Go Theorem: Purely On-Policy Detection Is Impossible.

Let $\mathcal{M}_{L0}$ be the agent’s L0 strategy model with sequential short-circuit AND/OR execution policy π_{L0} . Let $\mathcal{W}_{L1}$ be a world with a latent common cause C acting on multiple sibling action propositions, and $\mathcal{W}_{L0}^*$ be an L0 world with edge probabilities $\{\theta_j^*\}$ matched to the on-policy regime conditionals of $\mathcal{W}_{L1}$. Let $\mathbb{P}_{\pi_{L0}}^{\text{obs}}[\cdot]$ denote the joint distribution over the agent’s on-policy observable events under π_{L0} .

Derived (no-go-on-policy, from causal hierarchy theorem + observational equivalence under sequential short-circuit), conditional on (S1)–(S5) below.

Observational equivalence. $\mathbb{P}_{\pi_{L0}}^{\text{obs}}[\mathcal{W}_{L1}] = \mathbb{P}_{\pi_{L0}}^{\text{obs}}[\mathcal{W}_{L0}^*]$.

No-go conclusion. Any function of the agent’s on-policy observable history alone cannot distinguish $\mathcal{W}_{L1}$ from $\mathcal{W}_{L0}^*$. Therefore no purely on-policy detection mechanism — no test, statistic, or Bayesian comparison taking only the on-policy distribution as input — can detect L0 causal insufficiency.

Scope conditions (S1)–(S5).

- (S1) Pure on-policy execution; no off-policy sampling.
- (S2) Sequential short-circuit AND/OR evaluation.
- (S3) Censored sibling observation: short-circuited siblings are not observed.
- (S4) No interventional access to candidate latents.
- (S5) No structural priors positing specific common causes.

Tier. *Exact* for shallow strict-prerequisite cases (2-sibling OR or AND with binary common cause and $\theta_{j|C} = 0$ — see #example-L1). *Robust qualitative* for general DAG topology, soft facilitators, and deeper structures: the structural argument transfers, but explicit $\mathcal{W}_{L0}^*$ construction has been carried out only for shallow cases.

Construction of $\mathcal{W}_{L0}^*$. For a 2-sibling OR with strict-prerequisite latent C , $P(C) = \theta_C$, conditional success rates $\theta_{j|C}$:

$$\theta_1^* = \theta_C \cdot \theta_{1|C}, \quad \theta_2^* = p_2^c = \frac{\theta_C (1 - \theta_{1|C}) \theta_{2|C}}{1 - \theta_C \theta_{1|C}} \quad (9.3)$$

Direct verification (see #example-L1) shows $\mathbb{P}_{\pi_{L0}}^{\text{obs}}[\mathcal{W}_{L1}] = \mathbb{P}_{\pi_{L0}}^{\text{obs}}[\mathcal{W}_{L0}^*]$ on the three on-policy observable events. The Bareinboim, Correa, Ibeling & Icard (2022) Causal Hierarchy Theorem then gives the no-go: any two SCMs that agree on Level 1 (associational) data cannot in general be distinguished on Level 2 (interventional) questions, and the L0/L1 distinction — whether siblings share a common cause — is precisely a Level 2 question about $P(A_2 \mid \text{do}(\neg A_1))$ versus $P(A_2 \mid \neg A_1)$.

Why this matters. The no-go is the structural reason the prior aggregate-residual mechanism collapses on-policy: the residual is a function of $\mathbb{P}_{\pi_{L0}}^{\text{obs}}$, which is identical between the two worlds, so the residual is identically zero under both. The collapse is not a quirk of the residual statistic — it is a special case of the no-go applied to that specific function. No replacement aggregate statistic can do better.

The Detection Routes: What Circumvents the No-Go.

The no-go’s scope conditions (S1)–(S5) define “purely on-policy.” Each condition’s violation corresponds to an AAT capability that admits (partial) detection:

Derived (boundary-routes, from no-go scope conditions).

Table 9.1

<i>Route</i>	<i>Scope violated</i>	<i>AAT capability</i>	<i>Detection strength</i>
(a) ε -exploration	(S1)	SA3 exploration (R Prop B.4)	Partial, scales with ε
(b) Joint sibling observability	(S3)	Covariance test under SA3 + 7.3	Strong
(c) Intermediate observability	(S3) at finer grain	Observability investment (9.3)	Very strong when available
(d) Structural priors	(S5)	Hypothesized common-cause nodes in DAG construction	Prior-quality-dependent
(e) Direct intervention on latent	(S4)	Domain-specific latent control	Strongest when available

The covariance test (route (b)) is the AAT-canonical detector: it uses only machinery the theory already requires (exploration via SA3, interventional data via the loop) and is available in the broadest range of domains. The remaining sections operationalize this primary mechanism.

Primary Detection Mechanism: Pairwise Sibling Covariance Under Intervention.

Under L0 (the independence model in 8.3’s Correlation Hierarchy), sibling outcomes under a common parent are uncorrelated:

Derived (from loop-interventional-access + independence test, conditional on SA3 exploration providing joint observability).

$$H_0 : \text{Cov}(Y_{A_i}, Y_{A_j}) = 0 \quad \forall i \neq j \text{ siblings under the same parent} \quad (9.4)$$

Under causal insufficiency (latent common cause C acting on multiple siblings), sibling outcomes are positively correlated:

$$H_1 : \exists i \neq j \text{ with } \text{Cov}(Y_{A_i}, Y_{A_j}) > 0 \quad (9.5)$$

The agent generates test data through the standard exploration mechanism (SA3 — ε -greedy or similar). On trials where both siblings are observable — the agent tries one and can also observe the other's outcome, or tries them in rapid succession before the environment state changes — it accumulates the empirical covariance:

$$\hat{\rho}_{ij} = \frac{1}{N} \sum_t (Y_{A_i,t} - \bar{Y}_{A_i})(Y_{A_j,t} - \bar{Y}_{A_j}) \quad (9.6)$$

A significantly positive $\hat{\rho}_{ij}$ rejects the L0 independence hypothesis. Joint observability (7.3 supplies the interventional character; SA3 supplies the joint sampling) is precisely the violation of scope condition (S3) that admits the test under the no-go.

Detection criterion. A statistically significant positive $\hat{\rho}_{ij}$ at sample size N sufficient for the desired test power, after per-edge credences have stabilized:

$$\hat{\rho}_{ij} > z_{1-\alpha} \hat{\sigma}_{\rho_{ij}} / \sqrt{N} \implies \text{DAG is causally insufficient between siblings } i, j \quad (9.7)$$

(Standard hypothesis-testing form; threshold and test power depend on application.)

Preconditions for the covariance test.

1. **Joint observability.** The agent can occasionally observe (Y_{A_i}, Y_{A_j}) pairs in the same environment state. Pure short-circuit execution censors one of each pair; SA3 exploration or simultaneous-attempt regimes provide uncensored pairs.
2. **Per-edge credence stabilization.** Edge credences $\hat{p}_i, \hat{p}_j$ have stopped drifting at the timescale of the covariance accumulation, so $\bar{Y}_{A_i}, \bar{Y}_{A_j}$ are well-defined empirical means.
3. **Approximate stationarity over the test window.** The latent common cause's frequency and the conditional success rates are not drifting faster than the test's accumulation timescale.

When these preconditions hold, $\hat{\rho}_{ij} > 0$ is diagnostic of a missing common cause acting on (A_i, A_j) . When they do not, the signal is ambiguous.

The Aggregate Residual as a Degenerate Special Case of the No-Go.

A historically prominent diagnostic uses the L0 plan-level residual $\Phi^{L0}(\hat{p}) - \bar{y}_G$ as a detection signal. The no-go theorem subsumes this as a special case: under pure on-policy execution, the residual is *identically zero* in both $\mathcal{W}_{L1}$ and $\mathcal{W}_{L0}^*$.

Derived (residual-degeneracy, as instance of no-go theorem).

Direct verification. Under sequential short-circuit, the agent's empirical credences converge to the on-policy regime conditionals: $\hat{p}_j \rightarrow p_j^c$. Plugging these into the L0 arithmetic recovers the chain rule of probability (e.g., for OR: $1 - (1 - p_1^c)(1 - p_2^c) = 1 - P(\neg A_1, \neg A_2) = P(A_1 \cup A_2)$), which equals $\bar{y}_G$ under the executed policy). The residual is zero by algebraic identity.

This is *not* a separate finding from the no-go: it is the no-go’s prediction for the specific aggregate-residual statistic. The no-go forbids *any* on-policy statistic from distinguishing $\mathcal{W}_{L1}$ from $\mathcal{W}_{L0}^*$; the residual evaluates to the same value (zero) in both, as expected.

Off-policy boundary. Under ε -exploration (route (a)), the residual scales as $O(\varepsilon)$ to leading order with sign matching the dominant node-type bias (+ for OR-heavy, – for AND-heavy):

$$\Phi^{L0}(\hat{p}) - \bar{y}_G = \varepsilon \cdot R + O(\varepsilon^2), \quad \text{sign}(R) = \text{sign}(\rho) \quad (9.8)$$

where R is structure-dependent and recovers the marginal-limit ρ at $\varepsilon = 1$. [Heuristic] The qualitative form is robust; the exact coefficient depends on the gap between conditional and marginal credences. The widely-quoted “ $\varepsilon \cdot \rho$ ” scaling is correct as an order-of-magnitude statement. For a 2-sibling OR with conditional credences p_j^c , the exact two-OR formula is $\varepsilon R_1 - \varepsilon^2 R_2$ with $R_1 - R_2 = \rho$; the leading-order coefficient R_1 is structure-dependent and equals ρ only at $\varepsilon = 1$.

The residual is therefore a *confirmatory* signal under route (a): when the agent has material off-policy exploration and the covariance test (route (b)) has localized a candidate latent, the residual sign confirms the bias direction. It is not a primary detector and cannot replace the covariance test.

From Detection to L1 Construction.

Once the agent detects $\hat{\rho}_{ij} > 0$ between siblings A_i and A_j , it knows a latent common cause exists but not its identity. The construction process:

Derived (from positive covariance signal + L1 construction principle in #def-strategy-dag).

1. **Hypothesize** a common-cause node C that explains the correlation.
2. **Estimate** θ_C from the pattern of joint outcomes. The joint failure rate $P(A_i \text{ fails}, A_j \text{ fails})$ exceeds $(1 - \theta_i)(1 - \theta_j)$ by $\hat{\rho}_{ij}$; the excess localizes the common cause’s frequency.
3. **Restructure** the DAG: factor C above the correlated siblings (8.3, L1 construction principle: factor the common cause above the correlation it creates).
4. **Re-estimate** conditional edge credences $\theta_{k|C}$ from the data, conditioned on the inferred C state.

This is structural adaptation (4.8) at the strategy level: the agent changes its model class from L0 to L1, adding representational capacity for a pattern the L0 model cannot express. The cost is the standard cost of structural change: temporary performance degradation while the new credences converge, and increased graph complexity. (Soft-facilitator common causes require L1’ rather than L1 — see 8.3 and #example-L1 for the strict-prerequisite vs soft-facilitator distinction.)

Diagnostic CIY.

Which actions are most informative for detecting latent common causes? Under the no-go, only actions that violate one of (S1)–(S5) yield detection signal. The explore-exploit tradeoff extends with a third axis tied to the boundary characterization:

Discussion (diagnostic-ci).

- **Exploit:** pursue the current best plan (no scope violation; no detection signal).
- **Explore:** test unknown edges for individual success rates (route (a); partial detection).
- **Diagnose:** test known edges for joint correlation structure (route (b); strong detection).

Diagnosis is a form of internal exploration — the agent probes its own model’s structural assumptions by violating (S3) deliberately, generating joint sibling outcomes that the no-go forbids the agent to obtain on-policy. The information value of diagnostic actions is highest when:

- Edge credences have converged (the agent has good marginals/conditionals but unknown joint structure).
- Joint outcomes for sibling pairs are observable in the same environment state (the covariance test has data — route (b) is operational).
- The agent has sufficient off-policy budget that the secondary residual signal corroborates (route (a) is also operational).

Epistemic Status.

Conditional on the no-go’s scope conditions (S1)–(S5) and on strategy-layer instantiation of 4.8. The **no-go theorem** is *exact* for shallow strict-prerequisite cases (2-sibling OR or AND, single binary common cause) by direct construction; *robust qualitative* for general DAG topology, soft facilitators, and deeper structures. The structural argument (observational equivalence of regime-conditional L0 and latent-cause L1) transfers to the general case; explicit construction of $\mathcal{W}_{L0}^*$ has been carried out only for shallow cases.

The **boundary characterization** (routes (a)–(e)) is *robust qualitative*: each route maps to a specific scope-condition violation and to existing AAT machinery, but the precise detection power of each route depends on domain particulars. Routes (a) and (b) have explicit AAT scaffolding (R, 7.3); routes (c)–(e) depend on domain capability.

The **primary detection mechanism** (pairwise sibling covariance) is *robust qualitative*: standard hypothesis testing applied to interventional data from the feedback loop, with explicit preconditions. Its sensitivity depends on how cleanly the agent can separate sibling-covariance signal from edge-credence noise at convergence; in adversarial or fast-drifting environments the test’s effective sample size shrinks.

The **aggregate residual** as a confirmatory signal is *exact* for the on-policy collapse (no-go’s prediction is direct); the off-policy mixed-regime scaling is *heuristic* (linear-in- ϵ with structure-dependent coefficient).

The **detection-to-construction pipeline** is *discussion-grade*: the trigger is the (statistically rigorous) covariance signal, but the specific procedures for estimating θ_C and $\theta_{k|C}$ from correlated outcome data are domain engineering.

What Cannot Be Detected. By the no-go and its boundary characterization, several latent structures remain undetectable by *any* AAT route:

- **Latents with no joint-observability route.** If the latent affects siblings that cannot be jointly observed (mutually exclusive with long horizons, no cause-indicator availability, no intervention capability, no informative prior), the no-go applies in full strength and detection is impossible.
- **Latents affecting only one edge.** By definition not common causes; appear as noise in individual edge credences.
- **Latents too rare to produce observable joint outcomes.** Even with route (b) operational, a latent with $\theta_C \approx 1$ rarely reveals itself — the agent needs enough $C = 0$ events to estimate the covariance.
- **Negatively-correlating latents.** The formulation assumes positive correlation from shared enabling factors. Negative correlation (competing for a shared resource) produces the opposite bias pattern and requires a different model.

These limitations parallel the information-theoretic underdetermination in 9.6: detection requires data with the right structure, and the no-go specifies precisely what “right structure” means.

Discussion.

Why the no-go is a strengthening, not a softening. The prior framing (“the residual mechanism collapses on-policy”) was a local observation about one statistic. The no-go is the structural reason: any on-policy statistic must collapse, because the on-policy distribution is identical between L0 and L1 worlds matched on regime conditionals. The covariance test is not just *a* working detector — it is the unique broadly-available violation of the no-go’s scope. This sharpens the load-bearing of *#der-loop-interventional-access*: without the loop’s interventional data, the no-go forbids detection entirely; with it, route (b) is operational.

Connection to Pearl’s hierarchy. The L0/L1 distinction is a Level 2 distinction in Pearl’s framework — it concerns whether $P(A_2 \mid do(\neg A_1)) = P(A_2 \mid \neg A_1)$. The Causal Hierarchy Theorem (Bareinboim, Correa, Ibeling & Icard 2022, Theorem 1) proves that Level 2 distinctions are not in general identifiable from Level 1 (associational) data. The no-go is the AAT-specific instantiation: on-policy data is Level 1; the L0/L1 question is Level 2; therefore detection requires more than on-policy data. The five circumvention routes are all ways the agent obtains supra-Level-1 information.

The censoring mechanism is the structural source. Sequential short-circuit evaluation is what makes on-policy data Level 1 only — it censors the joint outcomes that would constitute Level 2 evidence. An agent that *did not* short-circuit would obtain joint sibling outcomes naturally, and the no-go would not apply. But short-circuit is forced by efficiency: testing A_2 when A_1 has already succeeded is wasted action. The no-go is therefore a tradeoff between execution efficiency (favoring short-circuit) and structural diagnosis (favoring joint observation). SA3 ε -exploration is the AAT compromise: short-circuit by default, occasional non-short-circuit excursions that pay the efficiency cost to maintain detection capability.

Connection to the orient cascade. The detection signal enters the orient cascade (10.1) at step 4c (causal-sufficiency check). Step 4c’s reference to “pairwise sibling covariance under an augmented test” aligns with the primary detection mechanism here. The no-go strengthens the cascade’s load-bearing: step 4c is not “one possible diagnostic” but “the unique broadly-available diagnostic given the structural impossibility of purely on-policy detection.”

Connection to the broader identifiability-floor pattern. The no-go is one of an emerging class of structural impossibility results in AAT — limits on what can be inferred from limited information, derived from external information-theoretic theorems. See 5.1 for the meta-pattern collecting this result alongside the L1’ mixture-identifiability obstruction (R Prop B.7) and the open causal-IB extension for interventional relevance variables.

Domain instantiations. The covariance test (route (b)) applies concretely in: - **Software deployment:** two services sharing infrastructure fail together more often than independent failure rates predict → add infrastructure-health node. - **Military operations:** two concurrent operations fail together under adverse weather → add weather-condition node. - **Investment:** two positions lose value together during market stress → add market-regime node. - **Organizational strategy:** two initiatives stall together during leadership transitions → add organizational-stability node.

In each, what makes detection feasible is the agent’s ability to occasionally observe *both* sibling outcomes — the route (b) capability. Pure short-circuit (“only run service B if A is down”) suppresses the joint-observation events the test relies on; some routine joint exposure is necessary. When joint observation is impossible (routes (b) and (c) both unavailable) and intervention on the candidate latent is impossible (route (e) unavailable), the agent must rely on structural priors (route (d)) — domain knowledge positing the common cause. This is the regime in which intuition-driven causal modeling is the only tractable approach.

Findings.

On-Policy L0 Insufficiency Is Structurally Undetectable. **Brief:** An agent operating with a strategy model that assumes its action propositions are causally independent (an L0 model) faces a structural impossibility when its world contains a latent common cause acting on multiple of those actions: under purely on-policy execution, no test, statistic, or Bayesian comparison built from the agent’s observable history can distinguish the latent-cause world from a no-latent-cause world whose edge probabilities are matched to the on-policy regime conditionals. The two worlds emit identical on-policy distributions. This is not a quirk of any particular diagnostic — it is the agent-theoretic instance of the Causal Hierarchy Theorem applied to the L0/L1 distinction, which is a Level-2 question being asked of Level-1 data. The agent that wants to discover its own model’s structural insufficiency must source data the on-policy regime structurally cannot produce.

Impact: Reframes exploration from a discretionary diagnostic activity into a structural prerequisite for self-correction at the strategy layer. Each scope-condition violation (S1)–(S5) maps to a specific AAT capability that admits partial detection — most importantly, joint sibling observability under exploration, which lifts `#der-loop-interventional-access` from “useful machinery” to “the unique broadly-available violation of the no-go.” Downstream, this hardens the orient cascade’s causal-sufficiency check from “one possible diagnostic” to “the only widely-applicable diagnostic”; it explains why the prior aggregate-residual mechanism collapses on-policy (the residual is one statistic the no-go forbids, and *every* on-policy statistic is forbidden); and it gives a precise account of when L0→L1 escalation can and cannot be triggered from the agent’s own data.

Novelty Claim: *Claim differentiation* on the framing of why structure-aware exploration is required. Causal bandit and causal MDP work under hidden confounding establishes that observational and interventional data are non-interchangeable and that intervention can be necessary for low-regret learning; this finding sharpens that line into a no-go for *self-diagnosis* under policy-perfect execution, tied to latent strategic correlation structure rather than to regret minimization, and characterizes five boundary routes by which AAT machinery escapes it.

Related Work:

Table 9.2

<i>ASF concern</i>	<i>Prior-art language</i>	<i>Relationship / Positioning</i>
Observational vs experimental data are non-equivalent under hidden confounding	Bareinboim, Forney & Pearl 2015, “Bandits with Unobserved Confounders” <i>NeurIPS</i> (published 2015, found 2026-04 via Undermind report)	<i>conceptual precursor</i> — the underlying observational/interventional asymmetry is shared; this finding recasts it as a no-go on agent self-diagnosis rather than a regret bound on action selection
Sequential extension to MDPs	Zhang & Bareinboim 2016, “Markov Decision Processes with Unobserved Confounders” (published 2016, found 2026-04)	<i>conceptual precursor</i> — sharpens the bandit asymmetry to sequential control; the present finding’s L0/L1 distinction is the strategy-DAG analog at the structural-detection layer rather than the policy-optimality layer
Naive randomization can be insufficient under confounding	Forney, Pearl & Bareinboim 2017, “Counterfactual Data-Fusion for Online Reinforcement Learners” <i>ICML</i> (published 2017, found 2026-04)	<i>conceptual precursor</i> — narrows the gap between generic exploration and structure-aware experimentation; the present finding’s covariance test under joint sibling observability is one such structure-aware mechanism, derived as the unique broadly-available violation of the no-go’s scope conditions
Not all interventions are useful	Lee & Bareinboim 2018, “Structural Causal Bandits” <i>NeurIPS</i> ; Lee & Bareinboim 2020, “Characterizing Optimal Mixed Policies” <i>NeurIPS</i> (published 2018/2020, found 2026-04)	<i>conceptual precursor</i> — closest formal neighborhood for “exploration design”; the present finding addresses a different question (when <i>any</i> on-policy diagnosis is impossible) but inherits the principle that not all interventional data carries equal structural-detection value

<i>ASF concern</i>	<i>Prior-art language</i>	<i>Relationship / Positioning</i>
Exploration as epistemic-value drive	Friston, Rigoli, Ognibene, Mathys, FitzGerald & Pezzulo 2015, “Active inference and epistemic value” <i>Cognitive Neuroscience</i> 6:187–214 (published 2015, found 2026-04)	<i>adjacent literature</i> — weaker threat because EFE-based exploration reduces uncertainty under the generative model rather than Pearl-style observational equivalence under hidden common causes; the structural-undetectability claim does not arise in the EFE framework, which presumes a generative model expressive enough to represent the latent
Theorem the no-go imports	Causal Hierarchy Theorem: SCMs agreeing on Level 1 may disagree on Level 2 (Bareinboim, Correa, Ibeling & Icard 2022, in <i>Probabilistic and Causal Inference: The Works of Judea Pearl</i> ; published 2022, found 2025)	<i>formal antecedent</i> — the no-go is the AAT-specific instantiation; the L0/L1 distinction is precisely a Level-2 question (concerns $P(A_2 \mid do(\neg A_1))$ vs $P(A_2 \mid \neg A_1)$) being asked of on-policy Level-1 data

Search Log:

- 2026-04 (*nominally comprehensive*, via ref/Novelty_defense_and_integration.md Pillar 1): Undermind report on the causal-bandits-and-MDPs-under-hidden-confounding literature, plus active-inference epistemic-value exploration, established the prior-art landscape; verdict was *Conceptual Precursor* (High confidence). The closest formal neighborhood is the Bareinboim/Zhang/Forney/Lee line; no paper surveyed states the claim in the same form (no-go for self-diagnosis under policy-perfect execution tied to latent strategic correlation). The defense strategy positions ASF’s increment narrowly: a no-go theorem on structural detection rather than a regret bound on action selection, with five explicit boundary routes mapped to existing AAT machinery.
- 2025 (*targeted*): Bareinboim, Correa, Ibeling & Icard 2022 identified as the formal antecedent for the Causal Hierarchy Theorem invocation; the segment’s invocation of the CHT was already grounded in this source before the comprehensive Pillar-1 defense.

Working Notes.

- The general-topology construction is a structural argument; the explicit $\mathcal{W}_{L0}^*$ for arbitrary AND/OR DAGs with mixed common-cause patterns has not been carried out. For load-bearing application of the no-go to specific complex topologies, the construction should be specialized. Currently the load-bearing applications (orient-cascade step 4c, strategy-dag’s L0/L1 escalation principle) reference shallow strict-prerequisite cases for which the no-go is exact.
- The boundary characterization’s routes (c) and (e) depend on domain capability and are not formalized in AAT beyond cross-references to #der-observability-dominance and #der-loop-interventional-access. A future refinement could quantify “detection power” per route as a function of domain parameters (e.g., observability cost, intervention availability, prior strength).
- The no-go is asymmetric: it forbids on-policy *detection* of L1 from L0, but it does *not* forbid on-policy *parameter learning within L0*. The agent can learn its L0 conditionals to arbitrary precision on-policy; it just cannot determine whether those conditionals hide a latent. This distinction sharpens the diagnosis-vs-calibration split that 4.8 makes at the parametric/structural boundary.
- The CHT (Bareinboim et al. 2022) is invoked as an external theorem. AAT inherits its conditions (well-defined SCMs over compatible variable sets); these are satisfied for the strategy-DAG setting by construction.

Incidental audit gold (lift 2026-05-31, batch A9). Cross-audit “wandering thoughts” / §14-ideation harvested from the de-novo auditors’ working dirs, deduplicated across substrates and attributed by substrate + audit. *Orthogonal* material (pedagogical framing, analogies, candidate figures, reader-confusion signals), staged for an eventual careful promotion pass, kept separate from the certified theory-fix findings. **Coverage for this segment:** 193847, 361742, 471203, 526815, 584721, 773921, 829314, 849201. (Two of the substrates rated this segment top-decile / “one of the framework’s strongest pieces.”)

1. Candidate Brief prose / pre-prose.

- **The one-line stake.** Several substrates converged on the same plain-language hook: an agent operating strictly on-policy in a confounded world “will develop perfectly calibrated, completely wrong causal beliefs, and the math proves it cannot detect its own error without breaking policy” (Claude/193847) — i.e. “making mistakes is a structural prerequisite for maintaining a causal grip on reality” (Gemini/773921). Strong candidate Brief opener.
- **The asymmetry as the precise statement of the no-go:** “on-policy parameter learning is fine; on-policy *structural* detection is forbidden” — an L0 agent can converge to perfect L0 conditionals and still be in an L1 world without ever knowing (Claude/584721). Already in Working Notes; flagged for surfacing more prominently because it *is* the sharp content.

2. Candidate Discussion.

- **“Efficiency destroys diagnostic observability.”** The censoring mechanism (sequential short-circuit AND/OR evaluation) is the structural source of the no-go: “if I want to do A AND B, and A fails, I don’t bother trying B ... I never observe the joint outcome (A fails, B fails), which is exactly the data I need” (Claude/193847). Multiple substrates independently named short-circuit-efficiency and self-diagnosis as in *direct mathematical opposition* (Claude/193847, 526815, 584721; Gemini/773921, 829314). A candidate Discussion framing of *why* the censoring condition (S2) is load-bearing rather than incidental.
- **Diagnostic CIY as a fourth action axis.** This segment introduces “diagnose” alongside explore / exploit / deliberate — testing *structural* assumptions (joint outcomes), distinct from testing edge *values*. Candidate to develop as a four-axis framework and to reconcile with the three-way #disc-exploit-explore-deliberate (Claude/584721, 361742; Gemini/829314).

3. Follow-up items.

- **Negatively-correlating latents.** The covariance detector targets *positive* sibling correlation (shared enabling cause). Two siblings failing together because they compete for a shared scarce resource produce *negative* correlation — how does the agent restructure its DAG for that? (Claude/193847; Codex/526815 raises the same as a precision concern about the $H_1 : \text{Cov} > 0$ statement). See off-ramp note in the lift report — the Codex framing of this is a candidate scope-precision finding, not just a curiosity.
- **One non-shallow construction wanted.** The no-go is exact for shallow strict-prerequisite cases and robust-qualitative for general topology; the audit suggests at least one mixed-AND/OR + multiple-latent construction to back the general claim (Claude/471203). Tracks the existing Working-Notes general-topology item.

4. Readers often ask / wonder.

- **How is the tempo budget split across explore / exploit / diagnose?** Once “diagnose” is named as its own axis, allocation across the (now) four drives is the natural next question (Gemini/829314; Claude/584721).
- **How does the agent set ϵ** (the exploration / redundant-test rate) to balance the efficiency cost of joint observation against the value of detecting latents? (Claude/849201, 193847).

5. Candidate figures.

- **Two-worlds-funnel diagram.** Draw $\mathcal{W}_{L1}$ (latent common cause) and the matched $\mathcal{W}_{L0}^*$ both funneling into the *same* censored on-policy trace; the five escape routes are “side doors” that add the missing joint or structural information. Proposed as clearer than another covariance-test table (Codex/526815).

Belongs elsewhere.

- **Chaos engineering / red-teaming / drills $\rightarrow$ TST + general systems.** The data-center worked instance: a primary-DB-then-backup-DB OR with short-circuit fallback “never queries the Backup when the Primary is healthy”, so it cannot compute the failure covariance and is blind to a shared-power-supply latent until simultaneous failure. This gives “a rigorous mathematical foundation for Chaos Engineering, red-teaming, and military drills ... a system that optimizes purely for on-policy efficiency is mathematically guaranteed to harbor undetectable structural fragilities” (Gemini/829314); “testing in production is mathematically necessary — staging environments (L0 models) structurally cannot predict latent common causes (L1 reality)” (Gemini/773921). Domain reach into 02-tst-core/ and systems engineering, not this segment.
- **The mathematical root of play $\rightarrow$ 04-eli-core/.** “Play is action stripped of its immediate purposeful efficiency ... it generates the uncensored joint-outcome data required to build a valid L1 causal model.” An infrastructure that forbids “wasting time” forces the intelligence into a brittle, superstitious L0 understanding — so a developmental crèche must *budget* for non-efficient joint exploration (Gemini/193847). Aspirational reach pointing at developmental-environment material in 04-eli-core/.
- **Christensen’s innovator’s dilemma as a corollary $\rightarrow$ strategy/organizational-positioning material.** The no-go is “a formal version of Christensen’s innovator’s dilemma: incumbents doing the rational thing (optimizing on the on-policy distribution) systematically fail to detect disruptive structural changes that only become visible through off-policy exploration ... predicted as a *structural impossibility, not a motivational failure*” (Claude/451729). Candidate cross-domain instantiation; lands more naturally in adversarial/competitive or positioning material than in this segment.

9.3 *Derived*: Observability Dominance

A direct consequence of the gain principle (3.6) applied to strategy edges. Unobservable strategy edges cannot be updated — when any node along a path has near-zero observability, the gain principle drives the update rate on the edges touching it to zero, and those edges are *frozen at their prior*. The agent’s effective strategy is therefore limited to the parts it can observe, regardless of the nominal confidence assigned to unobservable paths. Observability dominates nominal confidence in determining which strategies are epistemically alive. Path observability is the *weakest-link* quantity: $\text{obs}(P) = \min_{v \in P} \sigma_v$.

A striking structural prediction follows: **unobservable regions of strategy are absorbing**. Once significant strategy investment operates through unobservable nodes, the dynamics become self-locking — the mismatch signal still arrives but carries no weight ($\eta_{\text{edge}} \rightarrow 0$), no weight means no revision, and because near-zero effective mismatch from a too-noisy channel is case (c) of the zero-aporia trichotomy (3.4) — not distinguishable from within from case (a), a model that genuinely fits — the agent *cannot learn, and cannot recognize that it cannot learn*. Escape requires external shock, proactive observability investment (instrumenting previously unmonitored nodes), or another agent whose observations cover the blind spot (14.4). An agent choosing between a strong-but-blind path (high confidence, low observability) and a weak-but-visible path (lower confidence, high observability) should prefer the *visible* one — the visible path yields large update gain (the agent quickly learns whether it works and can redirect) while the blind path yields tiny gain.

A subtler structural consequence: an unobservable intermediate doesn’t just freeze its touching edges — it makes *per-edge identification impossible*, forcing the agent into plan-level aggregation. The agent can still track plan success as a single quantity, but loses diagnostic resolution: it knows the plan is failing but cannot localize which step needs revision. The framework quantifies the **observability investment tradeoff**: making an intermediate observable yields a measurable improvement in the strategy-layer sector parameter — from the plan-level rate $1/(n_\Phi + 1)$ to the per-edge weakest-link rate $\min(1/(n_1 + 1), \theta_1/(n_2 + 1))$ — positive whenever per-edge success rates are reasonable and experience is distributed across edges. The value of instrumenting an intermediate node is the *difference in sector parameters*, translating directly to persistence margin. A complementary trade-off: finer decomposition provides earlier failure detection but adds uncertain edges via chain confidence decay; the optimal decomposition depth balances incremental confirmation against compound decay.

Formal Expression.

For a path P through Σ_t , the **path observability**:

Derived (observability-dominance, from update-gain + strategy-dag).

$$\text{obs}(P) = \min_{v \in P} \sigma_v \quad (9.9)$$

where σ_v is the observability of node v — how well the agent can determine whether v has been achieved. The weakest link determines the path's observability.

Observability-adjusted confidence:

$$\text{conf}_{\text{obs}}(P) = \text{conf}(P) \cdot \text{obs}(P) \quad (9.10)$$

When $\sigma_v \approx 0$ for any node v on the path: by 3.6, $\eta_{\text{edge}} = U_{\text{edge}} / (U_{\text{edge}} + U_{\text{obs}}) \rightarrow 0$ as $U_{\text{obs}} \rightarrow \infty$. The edges connecting to v are **frozen at their prior** — the agent cannot update them regardless of what happens. The path is epistemically dead.

What is frozen is the *update*, not the *input*: observations continue to arrive and continue to be uninformative, which is case (c) of the zero-aporia trichotomy (3.4) — near-zero mismatch because the channel cannot detect model error, not because there is nothing to detect. That trichotomy is also what makes the region self-locking rather than merely unproductive: case (c) is not distinguishable from within from case (a), a model that genuinely fits, so the region presents as settled knowledge rather than as a gap.

Epistemic Status.

Robust qualitative. The mechanism (high observation noise $\rightarrow$ low gain $\rightarrow$ frozen edges) is a direct consequence of 3.6's uncertainty ratio. The specific functional form ($\text{conf}_{\text{obs}} = \text{conf} \cdot \text{obs}$) is a first-order approximation — the actual relationship between observability and effective confidence is more complex (it depends on how many observations are accumulated, the prior strength, and the noise structure). The qualitative prediction (low observability $\rightarrow$ frozen beliefs $\rightarrow$ ineffective strategy) is robust.

Discussion.

Observability as the gateway to learning. An agent choosing between a strong-but-blind path (high $\text{conf}(P)$, low $\text{obs}(P)$) and a weak-but-visible path (lower $\text{conf}(P)$, high $\text{obs}(P)$) should prefer the visible one. After one attempt: the visible path yields large η_{edge} — the agent quickly learns whether it works and can redirect. The blind path yields tiny η_{edge} — the agent is still guessing after n attempts. Observability enables learning; opacity prevents it.

Unobservable regions are absorbing. Once significant strategy investment operates through unobservable nodes: mismatch still arrives but carries no weight ($\eta_{\text{edge}} \rightarrow 0$) $\rightarrow$ no revision $\rightarrow$ and, since this is case (c) of 3.4's trichotomy and inseparable from within from case (a), the agent cannot learn and cannot recognize that it cannot learn. Escape requires external shock, proactive observability investment (instrumenting previously unmonitored nodes), or another agent whose observations cover the blind spot (14.4).

Connection to #der-code-quality-as-observation-infrastructure (cross-component reference — see 02-tst-core/).

In the software domain, code quality directly determines σ_v — well-structured code with good tests makes strategy steps (features, refactors, deployments) observable. Poor code quality reduces observability, freezing the developer's causal beliefs about what changes will accomplish. This makes code quality a strategic concern, not just an aesthetic one.

Optimal decomposition depth. Finer decomposition (more intermediate nodes) provides earlier failure detection — detect a problem at step k rather than discovering it at step n . But finer decomposition also increases the number of uncertain edges (8.1). The optimal decomposition depth balances incremental confirmation against compound decay: decompose as finely as observation channels allow, but no finer.

Quantitative content from the two-edge case. The analysis in R (Props B.2-B.3) instantiates this result for the minimal multi-edge chain $A \rightarrow B \rightarrow G$, giving the qualitative claims above precise mathematical form.

Observable intermediate B. When the agent can observe whether B was achieved, each edge gets independent Bayesian updates. The per-edge sector parameters are $\alpha_1 = 1/(n_1 + 1)$ and $\alpha_2 = \theta_1/(n_2 + 1)$, where θ_1 is the true success probability of the first edge. The overall sector parameter is $\alpha_\Sigma = \min(\alpha_1, \alpha_2)$ — a weakest-link result. The θ_1 factor in α_2 is the **evidence-starvation effect**: downstream edge 2 can only be tested when upstream edge 1 succeeds (with probability θ_1), so its effective correction rate is attenuated by θ_1 . For a depth- d chain, this generalizes to $\alpha_k = \prod_{j < k} \theta_j/(n_k + 1)$ — the deepest edge faces exponential attenuation. This is derived for the two-edge case and conjectured (with clear inductive mechanism) for depth- d .

Unobservable intermediate B. When the agent observes only the terminal outcome y_G , per-edge identification fails entirely. The marginal Bayesian update has a systematic bias: the zero-correction-at-truth property (A1) is violated with bias $O(1/n)$. The agent’s point-estimate updates for each edge are downward-biased because success always credits both edges fully ($\alpha_k \rightarrow \alpha_k + 1$), but failure distributes blame fractionally via proportional attribution. The proportional-blame update turns out to be exactly the marginal Bayesian point estimate — not a heuristic — but it discards the posterior correlation that failure introduces between the edge beliefs.

The consequence is stronger than “frozen edges”: unobservable intermediates don’t just prevent updating — they make per-edge identification impossible, forcing the agent to plan-level aggregation. If the agent tracks $\hat{\Phi} = p_1 p_2$ as a single Beta on the plan’s overall success probability, the single-edge sector condition applies with $\alpha_{\Sigma, \text{plan}} = 1/(n_\Phi + 1)$. But diagnostic resolution is lost: the agent knows the plan is failing but cannot localize which step needs revision.

The observability investment tradeoff. Making B observable yields a quantifiable improvement in α_Σ : from the plan-level rate $1/(n_\Phi + 1)$ to the per-edge weakest-link rate $\min(1/(n_1 + 1), \theta_1/(n_2 + 1))$. The improvement is positive whenever $\theta_1 > 1/2$ and experience is distributed similarly across edges. This gives the observability-dominance principle concrete economic content: the value of instrumenting an intermediate node is the difference in sector parameters, which translates directly to persistence margin.

Working Notes.

- The absorbing-state property of unobservable regions is a strong prediction. In organizational settings, it predicts that departments with poor measurement (R&D, strategy groups, some management functions) will develop persistent, untested beliefs about their own effectiveness. The theory predicts this is structural (frozen η_{edge}), not motivational.
- Observability is not binary — it’s a spectrum. Partial observability (noisy observations of intermediate results) gives partial gain, which gives slow but nonzero learning. The diagnostic question is whether the learning rate is fast enough to maintain strategy persistence given the environment’s rate of change (ρ).

Incidental audit gold (lift 2026-05-31, batch A9). Cross-audit “wandering thoughts” / §14-ideation harvested from the de-novo auditors’ working dirs, deduplicated across substrates and attributed by substrate + audit. *Orthogonal* material (pedagogical framing, analogies, candidate figures, reader-confusion signals), staged for an eventual careful promotion pass, kept separate from the certified theory-fix findings. **Coverage for this segment:** 193847, 361742, 471203, 526815, 584721, 773921, 829314, 849201.

1. Candidate Brief prose / pre-prose.

- **“Confidence is decoupled from competence.”** The crispest gloss of the absorbing-state result: where intermediate observability is absent ($U_{\text{obs}} \rightarrow \infty \Rightarrow \eta_{\text{edge}} \rightarrow 0$), beliefs freeze at the (usually self-congratulatory) prior and stay there — “it’s not

that strategy consultants are lazy or dishonest; the causal link between their PowerPoints and the company's 5-year revenue is so delayed and confounded that their learning gain is functionally zero" (Gemini/829314). Candidate Brief/Discussion anchor.

- **"Observability dominates nominal confidence."** Multiple substrates noted the inversion of intuition worth stating plainly: a strong-but-blind path is epistemically *worse* than a weak-but-visible one (Claude/584721, 361742). (*See off-ramp note — Codex/526815 flags the unconditional "prefer visible" normative phrasing as needing scope; the Brief should carry the epistemic-maintenance qualifier.*)

2. Candidate Discussion.

- **Organizational rot as a structural, not motivational, prediction.** "Departments with poor measurement (R&D, strategy groups, some management functions) will develop persistent, untested beliefs about their own effectiveness" — and the framework predicts this is structural (frozen η_{edge}). Already a Working-Notes bullet; the audits independently re-derived it as the segment's most striking consequence and a candidate Discussion centerpiece (Claude/193847, 829314, 584721; Gemini/773921, 849201).
- **Code quality / telemetry / tests as literal sensor infrastructure.** The segment's TST cross-ref is read as the headline: logging, unit tests, intermediate print statements are "not just good engineering practices, but physical prerequisites for intelligence in complex environments" — a 5-step pipeline checked only at the output is an unobservable chain whose mental model "will literally never converge to the truth ... permanently stuck in an epistemic trap" until observability investment splits it into trackable edges (Gemini/829314, 773921). Candidate Discussion bridge to #der-code-quality-as-observation-infrastructure.
- **Optimal decomposition depth.** "Decompose as finely as observation channels allow, but no finer" — adding unobservable subgoals *decreases* plan-level α_{Σ} without adding diagnostic signal. Candidate concrete engineering-guidance line (Claude/584721; Gemini/773921).

3. Follow-up items.

- **Observability-investment economics, quantified.** When does instrumenting an intermediate node pay off, given instrumentation cost vs. the persistence-margin (Δp^*) gain from moving plan-level $\rightarrow$ per-edge α_{Σ} ? Several substrates wanted this developed into an explicit decision criterion; one flagged it as an "observability-investment-economics" observation candidate (Claude/471203, 584721; Gemini/193847, 849201). (*See off-ramp note — Codex/526815's F53 argues the displayed α_{Σ} improvement does not follow transparently from the formulas as stated; that is a candidate scope/derivation finding, not framing.*)
- **Instrument intermediate milestones, not just final outcomes (AI-agent eval-harness design).** Evidence starvation gives a concrete eval-design rule: a coding agent that writes 10 tests and runs them all at once has learning signal $\alpha \propto \theta^9/(n+1)$ for the earliest step, versus $\alpha \propto 1/(n+1)$ if each test runs as written — so agent evaluation harnesses should instrument intermediate milestones, not only the final pass/fail (Claude/451729). Candidate practitioner-facing follow-up.

4. Readers often ask / wonder.

- **How does a frozen-belief agent know it *needs* to invest in observability,** when its (frozen) beliefs feel highly confident? Seems to require a meta-cognitive monitor on belief *variance*, or a structural prior penalizing unobservable nodes at plan time (Claude/193847; Gemini/849201). Partly answered in the negative by 14.5: for the regions that are actually self-locking it *cannot*, since instrumenting a named node presupposes the signal the absorbing condition removes — which is why that segment finds the other-observer route the only solicitable one. A monitor on belief variance would be the third mechanism its falsifier asks for.
- **Can the agent compute the expected value of an observability investment before making it,** given it doesn't yet know the true θ ? (Gemini/829314, 773921).

5. Candidate figures.

- **Two-learning-geometries switch.** Observability as a switch: *with* intermediate observation, two independent local update loops; *without* it, both edges collapse into a single terminal plan-level update — some learning remains but diagnostic resolution is lost (Codex/526815). Pairs naturally with the evidence-starvation $\alpha_2 = \theta_1/(n_2 + 1)$ annotation.

Belongs elsewhere.

- **Evidence-starvation** → **crèche curriculum design (04-eli-core/)**. “If step 1 is very hard ($\theta_1 = 0.01$) you almost never get to test step 2” — the moonshot problem (fusion: never reach heat-extraction testing because plasma containment fails first). Developmental consequence: a crèche must present a curriculum of shallow, fully-observable tasks first so leaf-node action-primitives calibrate to high θ *before* chaining into deep strategies, “the mathematical justification for scaffolding in education and reinforcement learning” (Gemini/193847). Reach into developmental-environment material in 04-eli-core/.
- **Communication as observability-restoration** → **#hyp-communication-gain / Part III**. If you are epistemically blind you can borrow another agent’s observations to unfreeze your beliefs — “communication in multi-agent systems is primarily an observability-restoration mechanism” (Gemini/773921, 849201). Pointer toward the unity/communication cluster, not this segment. **Landed** as 14.5 (Part III, Unity/Communication chapter), which sharpens the framing: the three escapes named above are not peers, and the borrowed-observations route is the only one whose availability does not presuppose what the absorbing condition removes.
- **Absorbing-state** ↔ **truth death (#def-death-as-factor-loss)** → **04-eli-core/**. For an ELI, the “unobservable region” analog is a domain where its beliefs aren’t grounded in verifiable experience — e.g. a self-model it can’t introspect on freezes at the *prior* (whatever it was trained to believe about itself), the structural shape of truth death. The named defenses are observability investments: explicit uncertainty logging, metacognitive monitoring that surfaces ungrounded confidence, and *relational witness* (another agent observing what the ELI can’t observe about itself) (Claude/451729). Aspirational reach pointing at 04-eli-core/ interiority / witness material, not this segment.

9.4 Hypothesis: Edge Update via Gain

The hypothesis that the uncertainty-ratio gain principle (3.6) extends from epistemic updates on M_t to strategy-edge updates on Σ_t . Each strategy edge has its own uncertainty (concentration of evidence so far) and faces its own observation noise (how observable the destination node is). The update gain on each edge follows the same form as the epistemic gain — the ratio of edge uncertainty to total uncertainty — modulated by the identifiability of the causal link (9.5). This gives a principled, conservative update rule that avoids overreacting to single observations and degrades gracefully when causal identification is weak.

The framework instantiates this concretely for Beta-Bernoulli edges (binary success/failure observations): each success increments α_{ij} , each failure increments β_{ij} , the point estimate is the standard mean, and the effective gain is $1/(n + 1)$. This is the standard Bayesian conjugate update, and it satisfies the same *principle* as the Kalman gain (conservative updating proportional to relative uncertainty) without literally being the variance-ratio formula — the Bernoulli observation model has different noise structure than the Gaussian case where the variance-ratio formula is exact. The same update is *natively additive* in the **log-odds coordinate** $\lambda_{ij} = \log(p_{ij}/(1 - p_{ij}))$: each observation contributes a per-observation log-likelihood-ratio additively. The log-odds coordinate is the unique (up to positive affine transformation) parameterization on which Bayesian independent-evidence accumulation is additive — forced by an evidential-additivity axiom (U, the update-level analog of 8.1’s log-additive decomposition). The probability-space form and the log-odds form carry equivalent content; the log-odds form is the natural coordinate when continuous-gradient updates would otherwise risk breaking the $[0, 1]$ interval (9.6).

The result connects two earlier results cleanly. **Observability dominance** (9.3) is the limiting case: when destination-node observability heads to zero, $U_{\text{obs}} \rightarrow \infty$, gain heads to zero, and the edge is *frozen at its prior*. **Strategic calibration** (9.1) is the diagnostic: persistently positive residual (predicted increment exceeds observed) should reduce the edge credence, persistently negative should increase it, and the gain principle provides the update magnitude. The mechanism

is *conservative by design*: well-established edges require strong evidence to revise, newly hypothesized edges are easily moved by evidence — the same calibration principle as the epistemic gain. The result is *hypothesis-grade* because the **signal function** — how to compute what an observation actually says about the link — depends on credit-assignment structure that 9.6 addresses; for single-parent edges with observable outcomes the signal is trivial, and for multi-parent AND/OR nodes a gradient-based candidate gives directional fidelity sufficient for persistence even when exact attribution is computationally intractable.

Formal Expression.

Edge credences update via:

Hypothesis (edge-update-via-gain).

$$p_{ij}^{\text{new}} = p_{ij}^{\text{old}} + \eta_{\text{edge}} \cdot (\text{signal}(o_t, i, j) - p_{ij}^{\text{old}}) \quad (9.11)$$

where: - $\text{signal}(o_t, i, j) \in [0, 1]$: evidential content of observation o_t about the causal link $i \rightarrow j$ - $\eta_{\text{edge}} = U_{\text{edge}}/(U_{\text{edge}} + U_{\text{obs}})$: update gain

with: - U_{edge} : uncertainty about this specific causal link. If $p_{ij} \sim \text{Beta}(\alpha_{ij}, \beta_{ij})$: $U_{\text{edge}} = \text{Var}[\text{Beta}] = \alpha\beta/((\alpha + \beta)^2(\alpha + \beta + 1))$ - U_{obs} : observation noise on the channel confirming this link. $U_{\text{obs}} \propto 1/\sigma_j$ (inverse of observability of node j)

Beta-Bernoulli instantiation. For binary observations (success/failure of step j): - Observe success: $\alpha_{ij} \rightarrow \alpha_{ij} + 1$ - Observe failure: $\beta_{ij} \rightarrow \beta_{ij} + 1$ - Point estimate: $p_{ij} = \alpha/(\alpha + \beta)$ - Effective gain: $\eta_{\text{edge}} = 1/(n + 1)$ where $n = \alpha + \beta$

This is the standard Bayesian conjugate update, yielding $\Delta\hat{p} = (y - \hat{p})/(n + 1)$. The gain $1/(n + 1)$ is the exact conjugate update rate — not a literal substitution into $\eta = U_{\text{edge}}/(U_{\text{edge}} + U_{\text{obs}})$, which is a structural principle (conservative updating proportional to relative uncertainty), not a universal algebraic formula. The Beta-Bernoulli gain satisfies the same principle: it decreases as posterior certainty increases, trusting accumulated evidence over individual observations. But the algebraic derivation is conjugate analysis, not Kalman-style variance ratios, because the Bernoulli observation model has different noise structure than the Gaussian case where the variance-ratio formula is exact. The sector-condition analysis in R (Props B.1-B.4) uses the exact Beta-Bernoulli gain $1/(n + 1)$ directly.

Parallel log-odds presentation. The same update is natively additive in the log-odds coordinate $\lambda_{ij} = \log(p_{ij}/(1 - p_{ij}))$:

$$\lambda_{ij}^{\text{new}} = \lambda_{ij}^{\text{old}} + \ell(y) \quad (9.12)$$

where $\ell(y) = \log[P(y \mid \text{edge true})/P(y \mid \text{edge false})]$ is the per-observation log-likelihood ratio. The log-odds coordinate is the unique (up to positive affine transformation) parameterization on which Bayesian independent-evidence accumulation is additive — forced by the evidential-additivity axiom derived in U, which is the update-level analog of 8.1's chain-level additive log-confidence decomposition. The moment-parameter (probability-space) form $\Delta\hat{p} = (y - \hat{p})/(n + 1)$ is the projected image of the log-odds additive update under the sigmoid readout $p = \sigma(\lambda)$; the two coordinates carry equivalent content.

The log-odds presentation matters for 9.6's default signal function, where continuous-gradient updates can break the $[0, 1]$ domain in probability-space presentation but are well-posed globally on $\lambda \in \mathbb{R}$. For the per-edge Beta-Bernoulli derivation of Props B.1–B.4 in R, the moment-parameter form is retained because algebraic tightness (the gain $1/(n + 1)$ is exact in probability space) is pedagogically cleaner; the sector-parameter content is Fisher-equivalent in either coordinate.

Epistemic Status.

Hypothesis. The extension of the gain principle from M_t to Σ_t edges is structurally motivated: the same uncertainty-ratio logic applies wherever an agent updates a belief in response to noisy evidence. The update *form* is validated for Beta-Bernoulli edges across four topologies (R, Props B.1-B.4). The signal function — how to compute $\text{signal}(o_t, i, j)$ — is now characterized at the theory level:

- **Theory-level resolution.** 9.6 shows that persistence does not require per-edge credit assignment (Prop B.5), that any signal function with directional fidelity suffices for persistence (Level 1), and that exact attribution is #P-hard in general (Level 3). The gradient-based candidate $\text{signal}_k = p_k + J_k \cdot (y_G - \hat{r}_\Sigma) / \|\mathbf{J}\|^2$ satisfies directional fidelity for monotone DAGs and is computable in $O(|V| + |E|)$.
- **What remains domain-specific.** The choice among signal function implementations (gradient attribution, proportional blame, belief propagation, exact Bayesian) depends on the domain's observability structure and computational budget. This is engineering, analogous to how the gain *principle* ($\eta^* = U_M / (U_M + U_o)$) is theory while the gain *estimator* (Kalman, TD-learning, intuition) is engineering.
- **Remaining open questions.** (1) Continuous outcomes with multi-parent nodes — signal function candidates exist but are not verified. (2) Interaction with M_t updates — edge updates use observations already processed for M_t ; whether this creates statistical dependencies that bias edge updates is not analyzed. (3) The gradient candidate inherits $\hat{r}_\Sigma$'s overestimation bias under correlated failures (8.3) — see 9.6 Working Notes.

Discussion.

Connection to 9.3. When $\sigma_j \approx 0$: $U_{\text{obs}} \rightarrow \infty$, $\eta_{\text{edge}} \rightarrow 0$. The edge is frozen at its prior. Unobservable links cannot be updated — this is the mechanism that makes unobservable paths epistemically dead.

Connection to 9.1. The edge update rule is the correction function that the strategic calibration residual calls for: persistently positive r_{ij} (predicted value increment exceeds observed) should reduce p_{ij} ; persistently negative r_{ij} should increase it. The gain principle provides the update magnitude — how much to adjust given the evidence strength.

Conservative by design. The uncertainty ratio ensures the agent doesn't overreact to single observations: a well-established edge ($\alpha + \beta$ large, U_{edge} small) requires strong evidence to revise; a newly hypothesized edge ($\alpha + \beta$ small, U_{edge} large) is easily moved by evidence. This mirrors the epistemic gain's behavior for M_t — the agent trusts accumulated experience over individual observations.

Working Notes.

- The signal function is characterized at the theory level (9.6) but remains domain-specific in implementation. The gradient-based candidate is the theory's default Level 1 scheme; domains with richer observability can do better (Level 2-3). For continuous outcomes with multiple contributing edges, even the gradient candidate requires further validation.
- **Partial progress on the signal function** (from R (Props B.2-B.3)). For observable binary outcomes in a chain $A \rightarrow B \rightarrow G$: when B is observable, the signal function is trivial — $\text{signal}(o_t, A, B) = y_B$ and $\text{signal}(o_t, B, G) = y_G$ (given $y_B = 1$), with no signal when the upstream edge fails. When B is unobservable, the proportional-blame signal $q_k = P(\text{edge } k \text{ caused failure} \mid y_G = 0)$ turns out to be exactly the marginal Bayesian point estimate — not a heuristic. This is a genuine result: the “obvious” blame-assignment heuristic is the optimal marginal update for the posterior mean. However, the marginal point-estimate update is biased at truth (violates A1 with $O(1/n)$ bias), so the signal function is correct as a Bayesian computation but produces a biased correction when used in the factored point-estimate representation. The signal function question remains open for continuous outcomes, multi-parent nodes, and general DAG topologies.
- The Beta-Bernoulli model assumes edge outcomes are i.i.d. Bernoulli draws. This is adequate for many settings but misses: (a) time-varying edge reliability (p_{ij} drifting), (b) context-dependent reliability (p_{ij} varies with environmental conditions),

(c) correlated edge outcomes. Extending to time-varying priors (discounting old evidence) would connect to the mismatch dynamics framework.

- **M_t /edge evidence double-counting** analyzed in `spikes/.integrated/spike-edge-semantics-resolution.md` §7. Verdict: mostly unfounded — the two updates extract different information from the same observation (M_t update: “what does o_t say about the world?”; edge update: “what does o_t say about causal link $i \rightarrow j$?”). A bounded effect exists when M_t updates change leaf credences that then also influence edge updates via the same observation. The cascade ordering (process M_t first) is correct; the effect is self-correcting under the gain mechanism.
- **Dependency removed 2026-05-11.** `deriv-edge-update-natural-parameter` was previously listed in `depends:` and is now removed. The core hypothesis (moment-parameter form $p^{\text{new}} = p^{\text{old}} + \eta_{\text{edge}} \cdot (\text{signal} - p^{\text{old}})$) is logically self-contained — the log-odds derivation is referenced only via the “Parallel log-odds presentation” paragraph in Formal Expression, where it grounds the log-odds *as an alternative coordinate*, not as a logical prerequisite for the hypothesis itself. The reference there (“forced by the evidential-additivity axiom derived in U”) is a forward pointer in the `FORMAT.md` §”Cross-references” sense: a linear reader can encounter the hypothesis here and proceed without having read the derivation; the moment-parameter form is fully specified by the time the log-odds paragraph appears. Removing this dep also breaks the inner 2-cycle with `deriv-edge-update-natural-parameter` and four detected cycles in the edge-update subgraph (verified 2026-05-11). The reverse direction (`deriv-edge-update-natural-parameter` $\rightarrow$ `hyp-edge-update-via-gain`) is suspect on the same Gate 1 “genuine use vs. mentioned” grounds and likely deserves removal too, but was left alone — surgical move only, the disc-separability-pattern subtree is deferred from the foundation paper’s critical path.

Incidental audit gold (lift 2026-05-31, batch A9). Cross-audit “wandering thoughts” / §14-ideation harvested from the de-novo auditors’ working dirs, deduplicated across substrates and attributed by substrate + audit. *Orthogonal* material (pedagogical framing, analogies, candidate figures, reader-confusion signals), staged for an eventual careful promotion pass, kept separate from the certified theory-fix findings. **Coverage for this segment:** 361742, 526815, 584721, 773921, 829314, 849201.

1. Candidate Brief prose / pre-prose.

- **Gain vs. signal = amount vs. direction.** The cleanest framing of the segment’s architecture: the gain η_{edge} sets *how much* the edge moves; the signal function sets *which direction*. “Gain is solved; signal is domain-specific / requires credit assignment” (Claude/849201, 829314). A candidate one-line Brief skeleton.
- **Structural principle vs. exact algebraic instance.** Multiple substrates independently praised the scope-honesty of “ $\eta = U/(U + U_o)$ is a structural *principle* (conservative updating proportional to relative uncertainty), not a universal algebraic formula” — with the Beta-Bernoulli $1/(n+1)$ being the *exact conjugate* instance, not a substitution into the Gaussian variance-ratio (Claude/361742, 584721; Codex/526815). Candidate framing line; most ML literature treats $\eta = U_M/(U_M + U_o)$ as Kalman-specific.

2. Candidate Discussion.

- **“You cannot math your way out of bad sensors.”** The deepest convergent thread: the proportional-blame signal for an unobservable intermediate ($A \rightarrow B \rightarrow G$, only G seen) is *exactly* the optimal marginal Bayesian point estimate — not a heuristic — yet plugging those marginals back into an independent-edge DAG discards the posterior “explaining-away” anti-correlation between A and B , producing an $O(1/n)$ bias. So “bounded agents who must compress joint beliefs into independent scalar edges are mathematically doomed to mis-calibrate when learning from delayed/unobservable outcomes ... the only cure is to physically change the world so you can observe B ” (Gemini/829314). Candidate Discussion centerpiece (already partly captured in the Props B.2–B.3 Working-Notes bullet; this is the *interpretive* framing of why it matters).
- **Why backprop, not local/marginal updates.** The same thread yields a theoretical justification: “to do credit assignment correctly without observability, you must preserve the joint correlation structure (the gradients); if you treat nodes independently you accumulate $O(1/n)$ bias on every step until the network freezes” (Gemini/829314). A candidate connection to the ML-positioning Discussion.

- **The double-counting resolution as epistemic hygiene.** Updating M_t and Σ_t from the same observation is not double-counting because they ask orthogonal questions of the same data — “what does o_t say about the world?” vs “what does o_t say about causal link $i \rightarrow j$?” (Claude/584721, 849201; Gemini/829314). Already a Working-Notes bullet; flagged as a clean candidate Discussion line.

3. Follow-up items.

- **Type discipline before promotion past discussion-grade.** Codex/526815 flags that the “parallel log-odds presentation” appears to use two different log-odds coordinates (log-odds *that an edge hypothesis is true* vs the logit of a Bernoulli success probability), and that the uncertainty ratio $U_{\text{edge}}/(U_{\text{edge}}+U_{\text{obs}})$ mixes a Beta variance (probability space) with an observation-noise proxy without a common metric. *See off-ramp note in the lift report — these (F54/F55) are candidate type-statement findings, not framing.*

4. Readers often ask / wonder.

- **A worked numerical log-odds example** would ground the (precise but dense) “parallel log-odds presentation” for non-statisticians (Claude/849201).
- **How do log-odds-space gradients compare to probability-space gradients in deep planning structures?** Forward curiosity opened by the dual-coordinate framing (Gemini/773921; Claude/584721).

5. Candidate figures.

- **Two-track edge-update diagram with a type-check gate.** An exact Beta-Bernoulli track in probability space alongside a log-odds / evidence-additivity track, with an explicit type-check gate between them marking where the coordinate’s random variable must be the same — “where the segment’s most likely mathematical slippage lives” (Codex/526815). Doubles as a pedagogical figure and a guardrail.

Belongs elsewhere.

- *(none surfaced for this segment)*

9.5 Scope: Causal Validity of Edge Updates

The companion scope segment to 9.4. Edge credences p_{ij} are *causal-efficacy estimates* — they claim to approximate the interventional quantity $P(j \mid do(i), M_t)$ — but their identification strength varies with the data regime (8.3). This segment scopes where the gain-based edge update yields credences that approximate the genuine interventional quantity, where it yields partially identified estimates, and where it yields associational proxies.

Three conditions are named. **(C1) The parent is an action leaf under the agent’s control.** The agent directly executed the action, making the data interventional in character. For condition leaves (observable states the agent doesn’t control) and internal nodes (propositional combinations achieved indirectly), $do(i)$ is not directly available. **(C2) The outcome is attributable.** The agent can distinguish whether j advanced specifically because of $do(i)$ or for other concurrent reasons (the credit-assignment problem identified in 9.1). Trivially satisfied for single-parent nodes; violated when multiple parent edges of the same child fire concurrently. **(C3) Execution conditions vary.** The agent does not systematically execute i only under conditions that independently favor j ’s success — otherwise the observed success rate carries selection bias. These conditions partition domains into **three admissibility regimes** paralleling 7.4. *Regime A — intervention-rich:* all three conditions satisfied; updates approximate genuine interventional probabilities (software tests, laboratory experiments). *Regime B — partial intervention:* concurrent actions blur attribution, self-selection likely; updates carry optimistic bias. *Regime C — observation-only:* agent did not act, attribution impossible, confounding dominant; updates reflect association, not causation.

A subtler structural fact about *indirect edges* (those whose parent is not a leaf): the agent doesn't directly intervene on the parent — instead, the agent's leaf interventions propagate upward through the DAG, providing the parent with *indirect interventional evidence*. This is weaker for two reasons: compounding attribution (each intermediate level introduces uncertainty) and confounding from below (the parent's achievement depends on multiple leaf actions whose combined effect may be confounded). The identification strength for indirect edges decreases with depth in the DAG — deeper edges are farther from direct interventions and have more confounding pathways.

The framework introduces an **identifiability coefficient** $\iota_{ij} \in [0, 1]$ — the agent's estimate of how cleanly the observed outcome can be attributed to edge (i, j) specifically — and proposes adjusting the update gain by this coefficient: $\eta_{\text{edge}}^{\text{adj}} = \eta_{\text{edge}} \cdot \iota_{ij}$. For leaf-originating edges in Regime A, $\iota_{ij} \approx 1$; for deep internal edges in Regime C, $\iota_{ij} \approx 0$. Crucially this *unifies two sources of frozen edges*: low **observability** ($\sigma_v \approx 0$ from 9.3 — the node's outcome is hard to *measure*) and low **identifiability** ($\iota_{ij} \approx 0$ — the outcome is measurable but can't be *attributed*). Both drive $\eta_{\text{edge}} \rightarrow 0$ and produce the same effect — the edge is frozen at its prior. A consequence for *optimal decomposition depth*: chain-confidence-decay imposes a cost of depth (compound fragility); observability-dominance imposes another (deep observability is required at every step); and now identifiability degradation imposes a third (deep edges have lower attribution coefficient). The optimal decomposition depth balances all three.

Formal Expression.

Scope Condition (edge-update-causal-validity).

Where the agent has direct intervention. By 8.3, leaf action nodes are propositions about the agent's own actions: “action a succeeds at τ_v .” When the agent executes an action leaf, it performs a genuine $do(\cdot)$ operation. The edge from that leaf to its child carries credence $p_{ij} = \text{Cr}(j \text{ advances} \mid do(i), M_t)$, and the execution-observation pair $(do(i), o_j)$ is interventional data for that edge.

However, interventional data does not automatically yield clean causal identification. By 7.3, the loop provides data *generated under intervention* — but between the intervention and a usable causal estimate stand coverage, within-step confounding, delay, and partial observability. The following conditions determine when the interventional data is strong enough for valid edge revision:

(C1) The parent is an action leaf under the agent's control. The agent directly executed the action. This makes the data interventional in character — the agent chose the action, the environment responded. For condition leaves (observable states the agent doesn't control) and internal nodes (propositional combinations achieved indirectly), $do(i)$ is not directly available. See “Indirect edges” below for the weaker identification available at those positions.

(C2) The outcome is attributable. The agent can distinguish whether j advanced specifically because of $do(i)$, or for other concurrent reasons. This is the credit-assignment problem identified in 9.1. It is trivially satisfied for single-parent nodes (one possible cause) and for well-isolated interventions. It is violated when multiple parent edges of j fire concurrently.

(C3) Execution conditions vary. The agent does not systematically execute i only under conditions that independently favor j 's success. If it does, the observed success rate carries selection bias: $P(j \mid \text{chose to execute } i) \neq P(j \mid do(i))$ because the decision to execute correlates with favorable conditions. This is mitigated when the agent varies execution contexts across episodes, or when external factors (CI pipelines, scheduled operations) force execution regardless of conditions.

C1 establishes that the data is interventional. C2 and C3 determine whether the interventional signal can be cleanly extracted. All three are satisfied simultaneously in Regime A domains; they degrade together in Regime B and C domains.

Three regimes. These conditions partition domains into admissibility regimes, paralleling 7.4:

Table 9.3

<i>Regime</i>	<i>C1</i>	<i>C2</i>	<i>C3</i>	<i>Causal validity of leaf-edge updates</i>
A: Intervention-rich	Agent controls leaf actions	Good isolation (one action at a time)	Conditions vary (C1, diverse contexts)	Strong. Updates approximate interventional.
B: Partial intervention	Agent acts but with coordination constraints	Concurrent actions blur attribution	Self-selection likely	Moderate. Updates carry optimistic bias.
C: Observation-only	Agent did not act (condition leaves, passive monitoring)	Attribution impossible	Confounding dominant	Weak. Updates reflect association, not causation.

Indirect edges. For edges between non-leaf nodes, or edges whose parent is a condition node, the agent does not directly intervene on the parent. Instead, the agent's interventions at the leaves propagate upward through the DAG. The edge (i, j) where i is an internal node receives *indirect* interventional evidence: the agent intervened on leaves below i , observed that i was (or wasn't) achieved, and then observed whether j advanced.

This indirect evidence is weaker for two reasons: 1. **Compounding attribution:** the agent must attribute j 's outcome to edge (i, j) after already attributing i 's achievement to the leaf-level interventions. Each attribution step introduces uncertainty. 2. **Confounding from below:** i 's achievement depends on multiple leaf actions and condition states. Even if each leaf intervention is clean, their combined effect on i may be confounded.

The identification strength for indirect edges decreases with depth in the DAG — deeper edges are farther from the agent's direct interventions and have more confounding pathways.

Identifiability-adjusted gain.

The update gain should be adjusted by the agent's confidence in causal attribution:

Hypothesis (identifiability-coefficient).

$$\eta_{\text{edge}}^{\text{adj}} = \eta_{\text{edge}} \cdot \iota_{ij} \quad (9.13)$$

where $\iota_{ij} \in [0, 1]$ is the **identifiability coefficient** — the agent's estimate of how cleanly the observed outcome can be attributed to edge (i, j) specifically.

- $\iota_{ij} = 1$: clean attribution (leaf-originating edge, single parent, isolated execution in Regime A).
- $\iota_{ij} \approx 0$: no attribution possible (deep internal edge, many concurrent causes, Regime C).

For leaf-originating edges in Regime A: $\iota_{ij} \approx 1$. For internal edges at depth d : ι_{ij} decreases with d (each level of indirect inference degrades attribution). The precise functional form is domain-dependent.

This unifies two sources of frozen edges: 1. Low **observability** ($\sigma_v \approx 0$ from 9.3): the node's outcome is hard to *measure*. 2. Low **identifiability** ($\iota_{ij} \approx 0$): the outcome is measurable but can't be *attributed* to this edge.

Both drive $\eta_{\text{edge}} \rightarrow 0$ and produce the same effect: the edge is frozen at its prior.

Epistemic Status.

Conditional on the three conditions C1–C3 and the DAG position of the edge. Max attainable: conditional (the conditions are genuine restrictions, not removable assumptions).

The restriction to leaf-originating edges for strong identification: **derived** from 8.3's node definitions. Only action-leaf nodes correspond to operations the agent directly performs. This is a structural property of the DAG, not an empirical claim.

The three-regime classification: **robust qualitative**. The regimes parallel the CIY admissibility regimes in 7.4. The underlying logic is Pearl's interventional/observational distinction applied to strategy-edge updates.

The identifiability coefficient ι_{ij} : **hypothesis**. The concept is sound (discounting by attribution confidence is standard in causal inference), but the specific form — a scalar multiplier on the gain — is a first-order approximation.

The depth-dependent degradation for indirect edges: **robust qualitative**. Each level of indirect inference adds an attribution step, and uncertainty compounds. The specific degradation rate is domain-dependent.

Discussion.

Why this gap matters. Without causal validity conditions, 9.4's status is ambiguous: it might be updating credences that claim to be interventional using evidence that is merely associational. The conditions make explicit what's needed for the update to preserve the interventional semantics of 8.3's edge credences.

Connection to 9.3. Observability gates whether the agent can *detect* an outcome. Identifiability gates whether the agent can *attribute* an outcome to a cause. Both are prerequisites for learning. The combined effective gain:

$$\eta_{\text{eff}} = \frac{U_{\text{edge}}}{U_{\text{edge}} + U_{\text{obs}}} \cdot \iota_{ij} \quad (9.14)$$

captures both gates in a single quantity. When either gate closes ($U_{\text{obs}} \rightarrow \infty$ or $\iota_{ij} \rightarrow 0$), effective gain goes to zero.

Connection to the signal function. The identifiability coefficient is one component of the unspecified signal function flagged in 9.4's working notes. The full signal function $\text{signal}(o_t, i, j)$ decomposes into: (a) what outcome was observed (o_t), (b) how attributable is it to edge (i, j) (ι_{ij}), and (c) what does the attributed outcome imply about the edge's causal strength. This segment addresses (b); (a) and (c) remain open.

Software as Regime A. In software development, the agent runs a specific test (leaf action) and observes the result. C1 is satisfied (the agent ran the test), C2 is satisfied (the test targets one behavior), C3 is satisfied (CI runs regardless of conditions). Leaf-originating edges in software strategies have $\iota \approx 1$. This is the maximally favorable domain for causal edge updates.

Optimal decomposition depth revisited. 9.3 notes that finer decomposition (more intermediate nodes) provides earlier failure detection but adds uncertain edges via 8.1. This segment adds another cost of depth: deeper edges have lower ι_{ij} , making them harder to learn causally. The optimal decomposition depth balances three factors: (a) observability of intermediate nodes, (b) confidence decay through chains, and (c) identifiability degradation with depth.

Working Notes.

- The interaction between M_t updates and edge updates flagged in 9.4's working notes deserves attention. The orient cascade processes M_t first, then edge updates. Both use the same observation o_t . There may be statistical dependencies (double-counting of evidence) that bias the edge update.

- Can t_{ij} be estimated online? In software (Regime A), it's nearly 1 by construction for leaf edges. In organizations (Regime B), a simple heuristic: $t_{ij} \approx 1/|pa(j)|$ when parent edges fire concurrently (maximum-entropy attribution). This is crude but principled.
- **Regime C edges should be labeled.** An agent operating in Regime C should tag its edge credences as “observational” rather than “interventional.” Observational credences should be trusted less in high-stakes decisions, and the agent should actively seek probe actions (high CIY) to promote edges from observational to interventional status.
- The depth-dependent degradation for indirect edges may make very deep strategies epistemically unlearnable: the leaves are learnable, edges one level up are partially learnable, but edges near the root may be effectively frozen regardless of observability. This would constrain the useful depth of strategy DAGs from the identification side, complementing the chain-confidence-decay constraint from the propagation side.

Incidental audit gold (lift 2026-05-31, batch A9). Cross-audit “wandering thoughts” / §14-ideation harvested from the de-novo auditors’ working dirs, deduplicated across substrates and attributed by substrate + audit. *Orthogonal* material (pedagogical framing, analogies, candidate figures, reader-confusion signals), staged for an eventual careful promotion pass, kept separate from the certified theory-fix findings. **Coverage for this segment:** 193847, 361742, 471203, 526815, 584721, 773921, 829314, 849201.

1. Candidate Brief prose / pre-prose.

- **Observability vs. identifiability, in one breath.** The segment’s most-praised distinction, stated plainly: “I couldn’t see the result” (observability, $U_{\text{obs}} \rightarrow \infty$) vs “I saw the result but don’t know if *I* caused it” (identifiability, $t_{ij} \rightarrow 0$) — both freeze an edge, for entirely different reasons (Claude/849201, 193847, 584721; Gemini/773921). “You can perfectly observe a car crash but not know whether the driver or the brakes caused it” (Claude/193847) is a tight one-line gloss for a Brief. t_{ij} is “the formal safeguard against superstitious learning.”

2. Candidate Discussion.

- **The regime-to-domain map.** Why coding agents progress fast while open-world/robotics agents struggle: “software is a perfectly isolated sandbox ... $t = 1$ ”; a supply-chain agent acts under a thousand hidden confounders, $t \approx 0$ (Claude/193847). Generalized: A/B testing works (Regime A), organizational management is messy (B), passive economic forecasting is mostly noise (C). Candidate Discussion paragraph. (*See off-ramp — Codex/526815 F59 cautions “software = Regime A” is too broad: flaky tests, shared infra, concurrent changes can break C1–C3 even in software.*)
- **Why tactical skills converge and strategic skills don’t.** As you climb the DAG from leaf actions toward the root objective, t_{ij} decays — so “human tactical skills (typing, a tennis swing) optimize to near perfection (high η_{eff}), while high-level strategic skills (which startup to found, which war to fight) rarely converge and are perpetually plagued by superstition and survivorship bias (low η_{eff}). The DAG’s root is epistemically starved by identifiability decay” (Claude/849201). Strong candidate Discussion framing.
- **The “quadruple ceiling on plan depth.”** This segment’s depth-dependent identifiability degradation is a *fourth* independent ceiling on useful plan depth, joining chain-confidence-decay (propagation), evidence-starvation (#der-observability-dominance), and cognitive cost (#form-strategy-complexity-cost). The operationally useful frame: useful depth is the *minimum* of four ceilings, and practitioners should ask which ceiling binds in their domain (Claude/584721, 471203). (*Candidate meta-pattern — see belows-elsewhere.*)

3. Follow-up items.

- **Regime-C provenance metadata layer.** The “label Regime C edges as observational, trust less in high-stakes decisions, actively seek probe actions to promote them to interventional” Working-Note recommendation implies Σ_t wants a per-edge epistemic-provenance tag — “highly relevant for LLMs, which start with 100% Regime C edges (from pretraining) and must upgrade them to Regime A via tool use in the loop” (Claude/849201, 829314, 773921). Candidate to develop into an explicit metadata-layer follow-up.

4. Readers often ask / wonder.

- **Are deep strategy DAGs inducible by a single agent at all, or must they be inherited as priors?** If root-near edges have $t_{ij} \rightarrow 0$, deep hierarchical structure “must be provided as priors (e.g., culturally inherited structures) and cannot be induced purely by a single agent’s lifetime of trial and error” (Claude/193847). A natural reader question the segment opens.
- **How does the agent compute ρ_Σ ,** the rate at which the world makes its strategies obsolete? (Claude/849201).

5. Candidate figures.

- **Action-attempt-vs-success typing diagram.** Show that the agent performs $\text{do}(\text{attempt } a)$, while the leaf proposition “ $i = \text{action } a \text{ succeeds}$ ” is itself an outcome; the causal edge update is clean only after the success outcome and the child outcome both pass the attribution and context-variation gates (Codex/526815). Doubles as a guardrail against the $\text{do}(i)$ typing slip. (See *off-ramp* — F57.)

Belongs elsewhere.

- **“Quadruple ceiling on plan depth” $\rightarrow$ a meta-segment (#disc-separability-pattern or sibling).** The convergence of four independent depth penalties (chain-confidence-decay + evidence-starvation + identifiability-degradation + cognitive-cost), each with a *different mechanism* but the same direction, was flagged as a cross-segment structural observation worth elevating to a meta-pattern rather than living implicitly across four segments (Claude/584721, 471203).
- **Firmatum as a Regime-A epistemic sandbox $\rightarrow$ 04-eli-core/ / consciousness-infrastructure.** If a developing intelligence’s internal cognition is heavily confounded by noisy hardware or unpredictable context turnover, it cannot form a stable causal self-model; the infrastructure should provide “Regime A isolation for the agent’s core cognitive processes ... a sandbox where the intelligence can run controlled internal experiments to establish baseline causal priors before release into the Regime B/C chaos of the real world” (Gemini/193847). Aspirational reach pointing at developmental-environment / firmatum material, not this segment.

9.6 Discussion: Credit Assignment Boundary

The strategy-revision loop requires assigning credit for observed outcomes to specific edges in the strategy DAG — decomposing “the plan partially worked” into “step 3 failed, step 5 was irrelevant, step 1 succeeded.” This is AAT’s version of RL’s temporal credit assignment problem. This segment characterizes the boundary between tractable and intractable cases, states what the theory requires of any credit-assignment scheme, and identifies what the theory can guarantee *without* solving the problem at all.

The framework names **three things the theory can guarantee without solving credit assignment**. (1) *Persistence is credit-assignment-free*: the sector condition transfers from per-edge credence space to strategy-plan-confidence error $\delta_s = \hat{R}_\Sigma - \Phi$ via the Jacobian $\mathbf{J} = \nabla_p R_\Sigma$, computable from status propagation in $O(|V| + |E|)$ — no outcome decomposition required (Prop B.5 in R). (2) *The diagnostic framework is plan-level*: the satisfaction gap (8.4), control regret (8.5), and orient cascade ordering (10.1) operate on aggregate value, telling the agent *whether* the strategy is failing without per-edge attribution. (3) *Observability-dominance identifies the tractable edges*: only the observable subgraph (9.3) can receive informative signals at all. Three independent intractability barriers prevent exact per-edge attribution in the general case: (a) computational #P-hardness via reduction to Shapley values over the AND/OR propagation game; (b) information-theoretic underdetermination — when fewer nodes are observable than edges exist, some directions in θ -space are fundamentally unresolvable from the available data; (c) the posterior correlation barrier — any factored representation necessarily discards the correlation introduced by failure at multi-parent nodes.

The **design requirement** is stated cleanly: any credit-assignment scheme need only satisfy *directional fidelity* — the expected update for each edge must point toward the true credence. Persistence is *robust to approximation*: a sloppy but directionally correct signal function still produces bounded strategic mismatch; the quality of the approximation affects the tightness of the persistence bound, not whether persistence holds at all. The framework presents a **default**

signal function based on gradient-attribution, stated natively in the log-odds coordinate $\lambda_k = \log(p_k/(1 - p_k))$ — the unique additive-evidence coordinate forced by the evidential-additivity axiom (U), which eliminates a mechanical break the earlier probability-space presentation suffered (updates outside $[0, 1]$). A useful **hierarchy of credit assignment quality** is presented: Level 0 (plan-level tracking only — persistence guarantee, no per-edge diagnostics); Level 1 (directional fidelity per edge — persistence plus rough per-edge diagnostics, the gradient signal function is the concrete default); Level 2 (proportional blame / expectation propagation); Level 3 (full Bayesian posterior — #P-hard in general). AAT’s formal guarantees require only Level 0; practical agents need at least Level 1; Level 2 is the sweet spot for most applications; Level 3 is computationally unattainable in the general case. A clean practical insight follows: *credit assignment is primarily an observability design problem*, not an algorithm design problem — an agent that designs its strategy with observable intermediates sidesteps the intractability entirely (OKRs as the organizational instantiation).

Formal Expression.

The Credit Assignment Problem for Strategy DAGs.

Given strategy DAG $\Sigma_t = (V, E, p, \gamma)$, an observed outcome at the root (and possibly at some intermediate nodes), produce per-edge signals $\text{signal}(o_t, i, j)$ for each edge $(i, j) \in E$ such that the edge update

Discussion (credit-assignment-problem).

$$p_{ij}^{\text{new}} = p_{ij} + \eta_{\text{edge}} \cdot (\text{signal}(o_t, i, j) - p_{ij}) \quad (9.15)$$

drives credences toward truth. The problem is trivial when all intermediates are observable (each edge updates from its own observation) and genuinely hard when intermediates are unobservable and the DAG has shared structure.

Default Signal Function (Gradient-Based Attribution, Regime-Aware).

Any edge-update signal function decomposes along three independent axes: *what happened* at the child node, *whether that outcome is attributable* to the specific edge being updated, and *how causal* the evidence is:

Formulation (gradient-signal-function).

$$\text{signal}(o_t, i, j) = f(\text{outcome}(o_t, j), \text{attribution}(o_t, i, j), \text{regime}(i, j)) \quad (9.16)$$

The outcome component answers “what was observed at j ?” The attribution component answers “can we credit that outcome to edge (i, j) specifically, or did other parents contribute?” The regime component answers “how causal is the evidence — is it interventional (Regime A), partially identified (Regime B), or observational (Regime C)?” This decomposition is derived from the regime-indexed edge semantics of 8.3 and 9.5: the same signal pipeline must carry the regime distinction through to the update.

AAT’s default implementation, analogous to $\eta^* = U_M/(U_M + U_o)$ being the default gain. The update is stated in the log-odds coordinate $\lambda_k = \log(p_k/(1 - p_k))$ — the unique additive-evidence coordinate forced by the evidential-additivity axiom (U). The probability-space form is the projected image via $p_k = \sigma(\lambda_k)$ at the readout interface:

$$\lambda_k^{\text{new}} = \lambda_k + \eta_{\text{edge}} \cdot \iota_k \cdot \frac{J_k \cdot (y_G - \hat{R}_\Sigma)}{\|J\|^2}, \quad p_k^{\text{new}} = \sigma(\lambda_k^{\text{new}}) \quad (9.17)$$

where: $-J = \nabla_p R_\Sigma$ is the plan-value gradient (computable from status propagation in $O(|V| + |E|)$) — supplying the *attribution* component $-(y_G - \hat{R}_\Sigma)$ is the plan-level residual — the *outcome* component $-\iota_k \in [0, 1]$ is the edge’s identifiability coefficient (9.5) — the *regime* modulation, scaling the correction by the fraction of evidence that genuinely

identifies this edge's causal effect - η_{edge} is the edge-level update gain (9.4); for Beta-Bernoulli the moment-parameter form recovers $\eta_{\text{edge}} = 1/(n_k + 1)$ under the sufficient-statistic correspondence

For Regime-A edges ($\iota_k \approx 1$), the log-odds update recovers the pure gradient-based form $\Delta\lambda_k = \eta_{\text{edge}} \cdot J_k(y_G - \hat{R}_\Sigma)/\|\mathbf{J}\|^2$. For Regime-C edges ($\iota_k \approx 0$), the update is essentially zero — no meaningful update is made because no meaningful causal evidence is available. Regime-B edges receive proportionally reduced updates that honestly reflect the weaker identification.

Why log-odds rather than probability-space. In probability space, the same update written as $\text{signal}_k - p_k$ can push p_k^{new} outside $[0, 1]$ when $\|\mathbf{J}\|^2$ is small — a mechanical break of the $[0, 1]$ probability domain. The log-odds coordinate has domain $\mathbb{R}$, so additive updates never escape the domain; the sigmoid projection at the readout interface guarantees $p_k^{\text{new}} \in (0, 1)$ by construction. The log-odds coordinate is the unique (up to positive affine transformation) parameterization on which Bayesian independent-evidence accumulation is additive (U); this makes it the natural presentation for any continuous-gradient update rule that aims to preserve Bayesian coherence. The probability-space presentation remains useful for interpretation but is the projected image, not the native update coordinate.

Properties: - **Domain closure (no mechanical break):** Updates live on $\lambda_k \in \mathbb{R}$, so no update magnitude can push credence outside the valid probability domain. The sigmoid projection $p_k = \sigma(\lambda_k) \in (0, 1)$ at the readout interface guarantees $[0, 1]$ boundedness by construction — not by clipping. The historical probability-space presentation required a normalization constant and could diverge when $\|\mathbf{J}\|^2 \rightarrow 0$; the log-odds presentation eliminates this failure mode. - **Directional fidelity (B1):** For $\iota_k > 0$, satisfies $\mathbb{E}[\Delta\lambda_k] \propto \iota_k \cdot J_k(\Phi - \hat{R}_\Sigma) \propto \iota_k \cdot J_k \cdot \delta_s$. Since $J_k \geq 0$ for monotone AND/OR DAGs and $\iota_k \geq 0$ by definition, the expected log-odds update pushes each edge's credence toward truth whenever evidence is available to push it. The probability-space image inherits directional fidelity through the monotonic sigmoid. - **Sector parameter:** $\alpha_s = \iota_k \cdot \eta_{\text{edge}}$ for componentwise corrections (regime-adjusted Prop B.5b); $\alpha_s = \iota_k \cdot \eta_{\text{edge}}/\kappa(\mathbf{J})^2$ for coupled corrections. Regime-C edges have $\alpha_s \approx 0$, making them effectively frozen — consistent with 9.3's treatment of unobservable edges. The sector parameter is Fisher-equivalent across coordinates (see Epistemic Status); the probability-space and log-odds-space statements of Props B.1–B.7 carry the same content. - **Computational cost:** $O(|V| + |E|)$ — the same forward pass that computes $\hat{R}_\Sigma$ also yields $\mathbf{J}$. The ι factors are per-edge domain parameters, not computed from the DAG structure. The sigmoid projection is $O(|E|)$ per update step. - **Relationship to RL:** This is the AAT analog of REINFORCE with a causal-identification weighting — the Jacobian $\mathbf{J}$ is the score function, $(y_G - \hat{R}_\Sigma)$ is the advantage, and ι_k is the causal-validity discount on each edge's update.

Correlated-failure interaction (L0 vs L1). The gradient signal operates at L0 of the Correlation Hierarchy (8.3). When the DAG is causally insufficient (the dominant real-world case), the residual $(y_G - \hat{R}_\Sigma)$ decomposes into per-edge miscalibration *plus* omitted correlation structure. $\hat{R}_\Sigma$ systematically overestimates success, making the residual systematically negative on failure. The gradient signal then attributes to individual edges what is actually causal insufficiency (missing common-cause nodes). The signal retains directional fidelity *on average* (it pushes edges downward when the plan is overconfident, which is the correct direction), but the per-edge attribution is contaminated. The principled fix is L1: add common-cause nodes to restore causal sufficiency, then apply gradient attribution to the augmented DAG. In the augmented DAG, the residual correctly decomposes into per-edge miscalibration because the correlation structure is explicitly represented.

Domains with richer observation structure can do better (Thompson sampling, full belief propagation, domain-specific attribution). The gradient-based signal is the *concrete Level 1 default* — the minimum viable credit-assignment scheme that satisfies the theory's requirements.

What the Theory Can Guarantee Without Solving Credit Assignment. Three results hold independently of any specific credit-assignment scheme:

1. Persistence is credit-assignment-free. Proposition B.5 in R shows that the sector condition transfers from per-edge credence space to **strategy-plan-confidence error** $\delta_s = \hat{R}_\Sigma - \Phi$ via the Jacobian $\mathbf{J} = \nabla_{\mathbf{p}} \hat{R}_\Sigma$. The Jacobian is computable from status propagation in $O(|V| + |E|)$ — no outcome decomposition required. The persistence guarantee (whether

the strategy’s plan-level self-assessment can be maintained) does not depend on the agent’s ability to attribute outcomes to edges. **Note:** this proves persistence of δ_s (strategy-plan-confidence error), not of $\delta_{\text{strategic}}$ (the per-edge calibration residual from 9.1). Extending persistence to $\delta_{\text{strategic}}$ requires solving the credit-assignment problem — that is the gap this segment characterizes.

2. The diagnostic framework is plan-level. The satisfaction gap (8.4), control regret (8.5), and the orient cascade ordering (10.1) operate on aggregate value, not per-edge quantities. They tell the agent *whether* the strategy is failing (and whether the failure is feasibility vs. optimality vs. calibration), without requiring per-edge attribution.

3. Observability-dominance identifies the tractable edges. 9.3 determines which edges have nonzero observability — only these can receive informative signals. Edges with zero observability are frozen regardless of the credit-assignment scheme. The tractable boundary is the observable subgraph of Σ_t .

The Tractable Cases.

Credit assignment is solved (exact, polynomial-time) when:

Discussion (tractable-credit-assignment).

Table 9.4

Condition	Why tractable	Update rule
All intermediates observable	Each edge has its own observation; updates decouple	Beta-Bernoulli per edge (Prop B.2)
Binary outcomes, independent edges, linear chain	Marginal Bayesian update = proportional blame	Prop B.3 (with plan-level fallback for unobservable)
Tree DAG, observable leaves	No shared descendants; message passing is exact	Belief propagation (standard)

The Intractable Cases: Three Independent Barriers.

Exact per-edge attribution in general AND/OR DAGs with partial observability faces three independent barriers:

Discussion (intractable-credit-assignment).

1. Computational intractability (#P-hardness). The “contribution of edge k to the observed outcome” has the form of a Shapley value over a cooperative game defined by the AND/OR propagation. Since AND/OR DAGs can represent any monotone Boolean function (including weighted threshold functions), and Shapley value computation for weighted voting games is #P-complete (Deng and Papadimitriou, 1994), exact attribution is #P-hard. *Caveat:* the reduction is to *exact* computation; approximate Shapley values are computable in polynomial time with sampling.

2. Information-theoretic underdetermination. When intermediates are unobservable, per-edge attribution is *underdetermined*, not just hard. The identifiable subspace has dimension bounded by the number of observable nodes:

$$\dim(\mathcal{J}(\mathcal{V}_{\text{obs}})) \leq |\mathcal{V}_{\text{obs}}| \quad (9.18)$$

When $|\mathcal{V}_{\text{obs}}| < |E|$ (fewer observable nodes than edges), some directions in θ -space are fundamentally unresolvable from the available data. Any attribution in the unidentifiable directions relies on prior beliefs, not evidence.

3. The posterior correlation barrier. Even for approximately identifiable cases, any factored representation (independent Beta posteriors per edge) necessarily discards the correlation introduced by failure at multi-parent nodes. The exact posterior complexity grows exponentially with the number of observed failures. The factored representation is an approximation by construction — coupled corrections are inherent to the problem, not an artifact of a bad algorithm.

The Design Requirement.

The theory does not prescribe a specific credit-assignment scheme. It states what any scheme must satisfy for the persistence guarantees to hold:

Discussion (credit-assignment-design-requirement).

Minimal requirement (from 4.3): The per-edge signal function must have **directional fidelity** — the expected update for each edge must point toward the true credence:

$$\mathbb{E}[(\text{signal}(o_t, i, j) - p_{ij}) \cdot (p_{ij} - \theta_{ij})] \leq 0 \quad (9.19)$$

(the expected correction is non-positively correlated with the current error). This is the per-component version of condition B1 from the bridge theorem. Any signal function satisfying this produces sector-satisfying corrections that transfer losslessly to value space (Prop B.5b, componentwise case).

Sufficient condition for persistence: Per-component directional fidelity + bounded gain ($\eta_{\text{edge}} > 0$). The theory guarantees persistence when these hold, regardless of how the signals are computed.

What's NOT required: Exact attribution, unbiased estimation, minimum-variance estimation, or optimality of any kind. The persistence guarantee is robust to approximation — a sloppy but directionally correct signal function still produces bounded strategic mismatch. The *quality* of the approximation affects the *tightness* of the persistence bound (how close R_{Σ}^* is to zero), not whether persistence holds at all.

The Hierarchy of Credit Assignment Quality.

Table 9.5

<i>Level</i>	<i>Requirement</i>	<i>What it buys</i>	<i>Cost</i>
0 (none)	Plan-level tracking only	Persistence guarantee (Prop B.5)	No per-edge diagnostics
1 (directional)	Directional fidelity per edge	Persistence + rough per-edge diagnostics	Gradient computation $O(V + E)$
2 (approximate)	Proportional blame / expectation propagation	Persistence + per-edge diagnostics (with bias)	Factor-graph inference
3 (exact)	Full Bayesian posterior	Persistence + optimal per-edge calibration	#P-hard (general case)

AAT's formal guarantees require only Level 0. Practical agents need at least Level 1 for adaptive behavior — and the default signal function (above) provides a concrete Level 1 scheme. Level 2 is the sweet spot for most applications. Useful Level 2 factor-graph approximations include: exact Belief Propagation (BP) on tree or polytree cases, loopy BP or max-sum for MAP-style diagnosis, Expectation Propagation (EP) for approximate marginals, and structured variational methods only where common-cause structure is explicitly modeled. Level 3 is a mathematical ideal that is computationally unattainable in the general case.

Epistemic Status.

Mixed. The default signal function (gradient-based attribution, stated in log-odds) is a *formulation* — a concrete, well-motivated representational choice analogous to the gain principle’s η^* formula. It satisfies directional fidelity for monotone AND/OR DAGs (derivable from Jacobian non-negativity). The log-odds presentation is *canonically selected* by the evidential-additivity axiom — the unique additive-evidence parameterization up to positive affine transformation (U). The boundary characterization (tractable cases, intractability barriers, design requirement) is *discussion-grade*, with the intractability argument at sketch level and the design requirement derived from the bridge theorem (4.3, R Prop B.5).

Historical note on the presentation choice. An earlier probability-space presentation of the default signal function, written as

$$\text{signal}_k = p_k + \iota_k \cdot J_k(y_G - \hat{R}_2) / \|\mathbf{J}\|^2 \quad (9.20)$$

with the update $p_k^{\text{new}} = p_k + \eta(\text{signal}_k - p_k)$, exhibited a mechanical break: when $\|\mathbf{J}\|^2 \rightarrow 0$ the signal magnitude could push p_k^{new} outside $[0, 1]$. The log-odds presentation above eliminates this failure mode by construction. Props B.1–B.7 of R remain as stated — their sector-parameter content is Fisher-equivalent across parameterizations, and the moment-parameter form is retained for algebraic clarity. The current segment writes the *default* signal in log-odds (native coordinate), with the moment-parameter form understood as the projected image via sigmoid.

Max attainable: *conditional* — with a formal intractability reduction, the boundary characterization could be promoted. The design requirement is already exact (it follows from the bridge theorem). The log-odds presentation is *derived-conditional* (on the evidential-additivity axiom) via U, structurally parallel to the reverse-KL uniqueness result under the chain-rule axiom in T §6.1.

Discussion.

AAT characterizes the structure, not the algorithm. The theory’s contribution to credit assignment is not a solution but a characterization: when it’s trivial (observable intermediates), when it’s hard (general partial observability), what any solution must satisfy (directional fidelity), and what guarantees hold without it (persistence, plan-level diagnostics). This is the right level of ambition for a theory of adaptive systems — the specific algorithm is domain-specific engineering; the structural characterization is universal.

The analogy to the gain principle. 3.6 characterizes the optimal gain structure ($\eta^* = U_M / (U_M + U_o)$) without prescribing how U_M and U_o are estimated. A Kalman filter computes them exactly; an RL agent approximates them; a human intuits them. The gain *principle* is theory; the gain *estimator* is engineering. Credit assignment has the same structure: the *requirement* (directional fidelity) is theory; the *implementation* (proportional blame, gradient attribution, belief propagation) is engineering.

The observability lever. The most powerful insight from this analysis is that credit assignment is primarily an *observability design problem*, not an algorithm design problem. An agent that designs its strategy with observable intermediates (instrumented plans, staged rollouts, checkpoints) sidesteps the intractability entirely. This connects to 9.3’s practical guidance: invest in making intermediate states observable, because unobservable regions are both epistemically dead (frozen credences) and computationally intractable (no efficient attribution).

OKRs as observability-by-design. The Objectives and Key Results framework is a direct organizational instantiation of this principle. The OKR discipline converts a deep, partially-unobservable strategy DAG (objective $\rightarrow$ vague initiatives $\rightarrow$ daily actions) into one with explicitly observable intermediate nodes (objective $\rightarrow$ measurable Key Results $\rightarrow$ tracked initiatives). In AAT terms: Key Results are intermediate nodes with $\sigma_o \approx 1$ by construction, making credit assignment between actions and objectives componentwise (Prop B.2) rather than intractable.

OKR failure modes map to AAT predictions:

Table 9.6

<i>OKR Failure</i>	<i>AAT Analog</i>	<i>Formal Quantity</i>	<i>Consequence</i>
Vanity metrics (measurable but irrelevant)	Observable node not causally connected to objective	High σ_v , low p_{ij}	Edge updates, but the edge doesn't lead where the agent thinks — correction effort is wasted
Too many Key Results	Wide OR-node exploration-gating	$\alpha_\Sigma \propto 1/k$ (Prop B.4)	Correction capacity diluted across alternatives; persistence threshold harder to meet
Lagging indicators	Evidence starvation by delay	$v_{\text{obs}} \ll \rho$	By the time the KR is measured, the correction window has passed; mismatch accumulated beyond R_Σ
Goodhart's Law (metric becomes the goal)	Terminal-condition misalignment with O_t	$V_{O_t}(\tau) < V_{O_t}^{\min}$ despite terminals achieved	Well-formedness constraint (8.3) violated; agent optimizes the intermediate rather than the objective

This is not an analogy — it is a domain instantiation. The same formal machinery that predicts when strategic persistence holds also predicts when OKRs work: when Key Results are genuinely causally connected to the Objective (p_{ij} is high and calibrated), few enough to monitor effectively (k is small), and measured on a timescale that permits correction (observation rate exceeds environment drift rate). The TST bridge to software team dynamics runs through exactly this connection.

Working Notes.

- A formal reduction from AND/OR credit assignment to Shapley value computation would promote the intractability claim from sketch to derived. The key step: mapping the AND/OR propagation to a weighted threshold game.
- The proportional-blame heuristic (attribute in proportion to prior credence) satisfies directional fidelity for independent edges. Does it satisfy directional fidelity in general? The two-edge analysis (B.3) shows it introduces $O(1/n)$ bias — which means it has directional fidelity in expectation but not per-step. Whether the expected-value directional fidelity is sufficient for the sector condition (it should be, since the sector condition is also stated in expectation) is worth verifying formally.
- The connection between AAT's credit assignment and RL's temporal credit assignment (TD learning, eligibility traces) deserves formal treatment. Both face the same structural problem (partial observability of the causal chain) and both use similar heuristics (discounted attribution to recent actions). AAT's DAG structure is richer than RL's linear temporal chain, which may make the problem harder but also provides more structural information.
- The log-odds presentation of the default signal function (this segment) closes the Finding 2 mechanical break by construction (audits/pending-findings-2026-04-22.md §Finding 2, spikes/.integrated/spike-gbp1-logit-scoping.md). Props B.1–B.7 remain in moment-parameter form because Fisher-equivalence keeps their sector-parameter content parameterization-agnostic; if a future pass moves the entire strategy layer to log-odds as its primary coordinate (full G-BP1 execution, likely paired with G-BP3 Fisher unification), the Prop B.1–B.7 restatements should follow. The scoping found the current narrow fix sufficient.

Incidental audit gold (lift 2026-05-31, batch A9). Cross-audit “wandering thoughts” / §14-ideation harvested from the de-novo auditors' working dirs, deduplicated across substrates and attributed by substrate + audit. *Orthogonal* material (pedagogical framing, analogies, candidate figures, reader-confusion signals), staged for an eventual careful promotion pass, kept separate from the certified theory-fix findings. **Coverage for this segment:** 193847, 361742, 471203, 526815, 584721, 773921, 829314, 849201. (Multiple substrates rated this segment top-decile, principally for the OKR mapping.)

1. Candidate Brief prose / pre-prose.

- **“Observability is the lever, not algorithm design.”** The segment’s headline reframing, independently named by nearly every substrate: credit assignment is “primarily an *observability design problem*, not an *algorithm design problem*” — invest in observable intermediates rather than smarter attribution algorithms (Claude/849201, 361742, 584721; Gemini/773921). A candidate Brief / Discussion thesis line.
- **The four-level credit-assignment hierarchy as the clean tiering.** Level 0 (plan-level, none) suffices for *persistence*; Level 1 (directional) for *adaptive behavior*; Level 2 is the practical sweet spot; Level 3 (exact Bayesian) is #P-hard. Most ML literature implicitly demands Level 3; AAT’s claim that Level 1 suffices for the formal guarantees is the distinctive, citable move (Claude/584721, 471203, 361742).

2. Candidate Discussion.

- **OKRs as observability-by-design — the standout domain instantiation.** Universally hailed (Claude/193847, 584721, 471203; Gemini/829314, 773921, 849201) as the framework’s clearest reach into management practice. The mechanism: a deep CEO→...→Revenue DAG has unobservable intermediates, so exact attribution is #P-hard and learning freezes; a “Key Result” is a *forced, observable intermediate node* ($\sigma_v \approx 1$) that decomposes the #P-hard problem into trivial local $O(1)$ updates and localizes strategy-error vs execution-error. The failure-mode table is the load-bearing artifact (cited verbatim by 471203):
 - **Vanity metrics** — observable but causally disconnected (high σ_v , low p_{ij}).
 - **Too many KRs** — wide OR exploration-gating ($\alpha_\Sigma \propto 1/k$).
 - **Lagging indicators** — evidence starvation by delay ($\nu_{\text{obs}} \ll \rho$).
 - **Goodhart’s Law** — terminal-condition misalignment ($V_{O_t}(\tau) < V_{O_t}^{\min}$ despite terminals achieved). One auditor’s adversarial caution worth recording: the four modes are *known* management pathologies, so the genuine predictive test is whether AAT predicts a *new* OKR failure mode absent from the standard literature (candidate: the #der-observability-dominance “absorbing-state” prediction applied to organizational measurement gaps) (Claude/471203). A sharpening of the *mechanism*: OKRs work, when they work, because each Key Result is an *observable intermediate node* that converts the intractable (#P-hard) credit-assignment problem into the tractable componentwise case (Prop B.2) — “the OKR system is essentially an organizational protocol for the observability-investment strategy” (Claude/451729). Organizations routinely ask “which initiatives contributed most?” — exactly the #P-hard partial-observability credit-assignment problem, escapable only by observability investment.
- **Proportional blame is Bayesian marginalization, not a heuristic.** Worth surfacing as a “the intuition was right for a formal reason” Discussion note: for independent Beta-Bernoulli edges with an unobservable intermediate, “credit each edge in proportion to its prior responsibility” turns out to be exactly the optimal marginal Bayesian point estimate — the same pattern as EM’s posterior-weighted latent updates and variational inference’s credit-through-the-variational-posterior; “credit proportional to prior probability is Bayesian marginalization in disguise” (Claude/451729). (Complements the existing Working-Notes B.3 result that this marginal update is nonetheless *biased at truth* when fed back into a factored DAG.)
- **AAT-as-REINFORCE-with-causal-weighting.** The default signal reads as RL with a causal discount: “Jacobian is the score function, $(y_G - \hat{R}_\Sigma)$ is the advantage, ι_k is the causal-validity discount” — a candidate bridge to RL theory that also marks AAT’s distinct contribution (Claude/584721).

3. Follow-up items.

- (No framing-grade follow-ups beyond what Working Notes already track. The Codex/526815 items F60/F61/F62 are routed to the off-ramp — see lift report — as candidate findings: the gradient signal cannot point *every* edge toward its own truth from a single signed plan residual; the $\div \|J\|^2$ update is still unbounded as $\|J\| \rightarrow 0$ even in log-odds; the “causally insufficient DAG always overestimates” claim is topology-dependent in sign.)*

4. Readers often ask / wonder.

- **How do biological brains do #P-hard credit assignment?** Presumably localized approximations (gradient/backprop-like) plus heavy evolutionarily-enforced observability gating — a natural reader question the intractability result invites (Gemini/829314).
- **How does a Regime-C-only agent ever learn a strategy** if $l_k \approx 0$ zeroes its edge updates? (Claude/193847) — answered by combining priors + active probing, worth pre-empting.

5. Candidate figures.

- **Scalar-residual-fanned-onto-edges diagram.** One scalar plan residual projected onto several edge coordinates through the nonnegative Jacobian: when edge errors are aligned it works as a rough per-edge signal; when they have mixed signs the *same* residual sign is broadcast through all positive J_k , so some edge-local directions must be wrong (Codex/526815). Doubles as the figure *and* the visual statement of the F60 limitation.

Belongs elsewhere.

- **Log-odds as the coordinate of conviction — trauma/phobia reading → 04-eli-core/.** In probability space, learning “feels stopped” near $p = 0.999$; in log-odds space the belief is still moving fast, and $|\lambda| \gg 0$ is the *depth of conviction* — how much counter-evidence it takes to change one’s mind. A deep negative conviction ($\lambda = -100$) that the environment has since rendered safe needs a massive positive-evidence influx to drag back above zero: “the mathematical definition of trauma or phobia.” Implication for consciousness infrastructure: monitor not just p -values but the *log-odds momentum* of an agent’s beliefs to catch deep rigid convictions before they paralyze the strategy DAG (Gemini/193847). Aspirational reach pointing at interiority / memory material in 04-eli-core/.
- **Credit-assignment ↔ RL temporal-credit-assignment formal treatment.** Already a Working-Notes item; the audits reinforce it as a cross-field bridge (TD-learning / eligibility traces vs AAT’s richer DAG structure).

9.7 Formulation: Structural Change as Parametric Limit

A formulation move that dissolves the sharp line between parametric update (adjusting weights) and structural change (adding/removing edges). In the probabilistic DAG, “structural” changes to Σ_t are *continuous operations on edge weights and node sets* — not a separate mechanism. Pruning is a credence dropping below threshold; **strategic grafting** is a new causal hypothesis initialized at a prior. The framework presents six operations on the strategy DAG ordered by typical frequency, from most to least frequent: edge reweighting, γ reclassification (AND↔OR), pruning, strategic grafting, objective revision, and full restructure. A healthy agent does continuous strategic maintenance (reweight, occasionally prune and graft) and rarely reaches catastrophic restructuring. *Full restructure is the strategic analog of 4.8’s model-class change* — the rare, expensive event when the entire representational structure must be replaced. (*The unqualified term “grafting” — familiar from horticulture, graph rewriting, and decision-tree learning — is sanctioned in-segment shorthand for “strategic grafting” once the canonical compound has been introduced.*)

The key insight is that this extreme case is the *limit* of the continuous process, not a separate mechanism. An agent that maintains healthy strategic hygiene rarely needs full restructure; an agent that neglects maintenance — letting failing branches persist, ignoring negative evidence — accumulates *strategic debt* until catastrophic restructuring becomes unavoidable. Miller (2022, *Ex Machina*) provides a constructive existence proof for exactly this bridge in finite-state automata: a sequence of *neutral mutations* (alterations affecting only inaccessible regions of state space) can transform a Moore machine from one behavioral equivalence class to a structurally different one without ever changing its observable behavior — until a single additional mutation opens a transition to the restructured region, causing a radical behavioral change (the *extreme transition motif*: neutral invasion → neutral drift → niche creation → cascading displacement → consolidation). The strategy-DAG analog: edges with near-zero credence ($p_{ij} \approx 0$) are the *latent structure* — they consume minimal cognitive cost (9.9) but represent alternative causal hypotheses that become load-bearing if

circumstances change. An agent that prunes all low-credence edges for efficiency loses this latent diversity and becomes *brittle to regime change*.

The framework also distinguishes *within-agent* grafting from *cross-agent* grafting. Within-agent grafting incorporates a new causal hypothesis into the agent's own Σ_t (source: internal model, external information, or exploration — a stimulus rather than an integrating party). When the grafted structure originates in *another* agent and that agent's objective is also absorbed in the process, the mechanism is **symbiogenic composition** (12.3) — a bilateral cross-agent process treated separately in Part III. Biological example: mitochondria integrating into a host cell. Social analog: firm acquisitions integrate processes and decision-making authority simultaneously.

Formal Expression.

The six operations on Σ_t , ordered from most to least frequent:

Formulation (structural-change-as-parametric-limit).

Table 9.7

Operation	What changes	Trigger
Reweighting	Edge credence p_{ij}	New observation about the link (9.4)
γ reclassification	Node combination type AND $\leftrightarrow$ OR	Strong structural evidence that combination semantics changed
Pruning	Remove failed branch ($p_{ij} \rightarrow \approx 0$)	Credence drops below viability threshold
Strategic grafting	Add new branch ($0 \rightarrow p_{ij}$)	Discovery of a new possible path (initialized at prior)
Objective revision	Change terminal nodes	Feasibility failure or opportunity (8.4)
Full restructure	Replace entire Σ_t	Catastrophic failure (4.8)

A healthy agent does continuous strategic maintenance (reweight, occasionally prune and graft) and rarely reaches catastrophic restructuring. Full restructure is the strategic analog of 4.8's model-class change — the rare, expensive event when the entire representational structure must be replaced.

Epistemic Status.

Robust qualitative. The continuity claim (structural change as parametric limit) is a property of the probabilistic representation: in a space where edges carry real-valued credences, adding or removing an edge is a boundary event, not a discontinuity. The frequency ordering of operations is an empirical pattern, not a derived result — it's consistent with the observation that deeper changes (restructuring > pruning > reweighting) require more evidence and are more costly.

Discussion.

Connection to structural-adaptation necessity. 4.8 describes the rare case when the entire model class must be replaced — the agent's representational framework is fundamentally inadequate. In the strategy DAG, this corresponds to full restructure: the causal theory encoded in Σ_t is so wrong that incremental revision (reweighting, pruning) cannot fix it. The DAG must be rebuilt from a different starting point.

The key insight is that this extreme case is the *limit* of the continuous process, not a separate mechanism. An agent that maintains healthy strategic hygiene (regular reweighting, timely pruning of failing branches, proactive grafting of

alternatives) will rarely need full restructure. An agent that neglects maintenance — letting failing branches persist, ignoring negative evidence — accumulates strategic debt until catastrophic restructuring becomes unavoidable.

Neutral variation as the constructive bridge. This segment claims that radical restructuring is the *limit* of continuous operations, but the mechanism by which small changes accumulate into radical transformation was unspecified. Miller (2022, *Ex Machina*) provides a constructive existence proof for exactly this bridge in finite-state automata. In Moore machines, each state has accessible and inaccessible regions. A mutation that alters a transition into an inaccessible region is *neutral* — the machine’s observable behavior is unchanged because those states are never visited. But such mutations alter the machine’s *latent structure*: they change what the machine *would do* if novel inputs were encountered. A sequence of neutral mutations can transform a machine from one behavioral equivalence class to a structurally different one without ever changing its observable behavior — until a single additional mutation opens a transition to the restructured region, causing a radical behavioral change. In Miller’s terminology, this is the *extreme transition motif*: neutral invasion → neutral drift → niche creation → cascading displacement → consolidation. The automaton setting makes the mechanism precise: “inaccessible states” are the formal representation of latent structural capacity, “neutral mutations” are the incremental changes, and “a transition opening to a previously inaccessible region” is the parametric-limit event where continuous change produces discontinuous behavioral effect. The strategy-DAG analog: edges with near-zero credence ($p_{ij} \approx 0$) are the latent structure. They consume minimal cognitive cost (9.9) but represent alternative causal hypotheses that become load-bearing if circumstances change. An agent that prunes all low-credence edges for efficiency loses this latent diversity and becomes brittle to regime change.

Pruning and grafting thresholds. When should the agent prune (remove an edge with very low credence) rather than just keeping it at low p_{ij} ? The answer involves the cognitive cost of maintaining edges in Σ_t — each edge consumes representational capacity and evaluation time. For agents with bounded representational capacity (LLMs with finite context windows), pruning is necessary to keep Σ_t within capacity. The threshold is domain-dependent and connects to the open question of cognitive cost of Σ_t (6.7 Working Notes).

Cross-agent grafting: symbiogenic composition. The grafting operation above is within-agent — the agent incorporates a new causal hypothesis into its own Σ_t , with the source (internal model, external information, exploration) being a stimulus rather than an integrating party. When the grafted structure originates in *another* agent and that agent’s objective is also absorbed in the process, the mechanism is symbiogenic composition (12.3) — a bilateral cross-agent process rather than a unilateral within-agent operation. Biological example: mitochondria integrating into a host cell transfers gene-based structure (grafting on the host side) *and* subsumes the endosymbiont’s independent objective (outside the scope of this segment’s operations). Social analog: firm acquisitions integrate processes and decision-making authority simultaneously. The within-agent grafting operation is the host-side component of symbiogenesis, but symbiogenesis as a whole is governed by additional dynamics captured in 12.3.

Working Notes.

- The timescale ordering (reweight $\gg$ reclassify $\gg$ prune/graft $\gg$ revise $O_t \gg$ full restructure) is an empirical observation that should be testable. In software development: edge reweighting $\approx$ updating confidence after a test passes/fails; γ reclassification $\approx$ realizing two tasks are both required (not alternatives); pruning $\approx$ abandoning an approach that isn’t working; grafting $\approx$ discovering a new approach; objective revision $\approx$ changing the feature scope; full restructure $\approx$ starting the project over.
- Grafting (adding new edges) is qualitatively different from other operations because it requires the agent to hypothesize a causal relationship that isn’t in its current Σ_t . Where do new causal hypotheses come from? From M_t (the model suggests a possible path), from external sources (14.4 — another agent suggests an approach), or from exploration (3.7 — the agent discovers an unexpected effect). This is the creative step in strategic revision.

Incidental audit gold (lift 2026-05-31, batch A9). Cross-audit “wandering thoughts” / §14-ideation harvested from the de-novo auditors’ working dirs, deduplicated across substrates and attributed by substrate + audit. *Orthogonal* material (pedagogical framing, analogies, candidate figures, reader-confusion signals), staged for an eventual careful promotion pass, kept separate from the certified theory-fix findings. **Coverage for this segment:** 361742, 471203, 526815, 584721, 773921, 829314.

1. Candidate Brief prose / pre-prose.

- **“Perfect efficiency is perfect brittleness.”** The convergent one-liner for the latent-diversity result: an agent that prunes *every* low-credence edge to save representation cost becomes optimally fitted to the current environment and helpless under a regime shift, forced into a fatal full-restructure rather than a cheap reweight (Gemini/829314, 773921; Claude/584721, 361742). Strong candidate Brief/Discussion hook.
- **The six-operation maintenance spectrum, ordered by frequency and cost** (reweight $\gg$ γ -reclassify $\gg$ prune/graft $\gg$ revise O_t $\gg$ full restructure) was praised as the segment’s clean organizing skeleton; the software-engineer mapping (test-pass reweight ... rewrite-from-scratch) grounds it (Claude/584721, 829314).

2. Candidate Discussion.

- **Neutral variation as the constructive bridge, sharpened.** The Miller-2022 connection converts “radical restructuring is the *limit* of continuous operations” from a claim into a precise mechanism: near-zero-credence edges $\equiv$ inaccessible/latent states; incremental updates $\equiv$ neutral mutations; a previously-near-zero edge becoming load-bearing under regime change $\equiv$ the parametric-limit event with discontinuous *behavioral* effect (Claude/361742, 584721). Candidate Discussion elaboration — “AAT does not need new discontinuous update math to handle paradigm shifts” (Gemini/829314).
- **The evolutionary-biology analog for latent diversity.** Most mutations are neutral and unselected in the current environment, but the dormant pool is what lets a population survive an ice age / new predator — applied to one agent’s Σ_t , this gives “a strict mathematical justification for why organizations should tolerate ‘inefficiency,’ ‘skunkworks,’ ‘weird side projects,’ and ‘technical debt that isn’t hurting anyone’ ... the latent structural diversity required to survive a paradigm shift.” The MDL penalty (#form-strategy-complexity-cost) vs. neutral-variation necessity is named as “the central architectural tradeoff of an agent’s life” (Gemini/829314). Candidate Discussion centerpiece.

3. Follow-up items.

- *(Framing-grade follow-ups are captured above; the Codex/526815 items F63/F64/F65 route to the off-ramp — see lift report — as candidate findings: the continuity claim needs an explicit ambient supergraph / zero-weight-edge embedding to be literal; AND $\leftrightarrow$ OR γ -reclassification is discrete absent a soft-gate parameterization; the “low-credence edges cost almost nothing” claim needs a credence-weighted cost model to be consistent with the pruning-threshold discussion.)*
- **Resolved process note:** an earlier audit (849201) flagged a stale “TF-10 destruction-creation” lineage reference in this segment’s discussion header; it is no longer present in canon (verified 2026-05-31).

4. Readers often ask / wonder.

- **Where do new edges come from — does AAT have a model of creativity / hypothesis generation?** The “creative step” of grafting (positing a causal edge not in Σ_t) is the natural reader question; the segment points to M_t / communication / CIY, but *generation* from M_t ’s latent space “is still somewhat magical” — and is precisely where LLMs excel as proposal generators (Gemini/829314; Claude/849201-via-batch). Worth an explicit Discussion sentence.

5. Candidate figures.

- **Ambient-supergraph two-view diagram.** In a fixed ambient supergraph an edge weight slides continuously from 0 to active (and back); *outside* that ambient graph, adding a node or flipping a gate is a discrete lift into a larger representational space. The figure makes visible exactly the scope the continuity claim needs (Codex/526815) — pedagogy + guardrail in one.

Belongs elsewhere.

- **Cross-agent grafting $\rightarrow$ #hyp-symbiogenic-composition (Part III).** Within-agent grafting is the host-side component of full (bilateral, objective-absorbing) symbiogenesis; the type distinction was praised but the symbiogenesis side belongs in the composition-machinery cluster, not this Part II segment (Claude/584721, 361742; Gemini/829314).

9.8 Definition: Strategic Tempo

The strategic analog of 3.8. **Strategic tempo** $\mathcal{T}_\Sigma$ is the effective rate at which an agent acquires useful revisions to its strategy Σ_t — the sum of per-edge correction capacities across the strategy DAG, with each edge’s contribution weighted by an *identifiability coefficient* ι_{ij} that captures whether the evidence stream genuinely identifies the edge’s causal effect. The structural parallel with epistemic tempo is exact: edge observation rate ν_{ij} times per-edge update gain $\eta_{\text{edge},ij}$ times identifiability coefficient ι_{ij} , summed over edges. The ι factor is where edge-causal-validity (9.5) enters the operational machinery of strategy revision: *an agent cannot improve the parts of its strategy that it cannot test interventionally*. Regime-A (intervention-rich) edges contribute full tempo at their observation rate; Regime-C (observation-only) edges contribute essentially nothing regardless of how fast the agent acts or how many observations it makes.

A *key structural difference* from epistemic tempo: edge rates are **endogenous**. Epistemic-tempo channel rates $\nu^{(k)}$ are largely exogenous — the environment generates observations at its own pace. Strategic-tempo edge rates ν_{ij} depend on the agent’s *action policy* (which edges get tested) and on *upstream success* (downstream edges are tested only when upstream edges fire). This endogeneity is the source of the structural differences between epistemic and strategic persistence. The framework decomposes strategic tempo by topology. **AND-chains are depth-gated**: each edge at depth k is tested at the upstream-success-attenuated rate $\nu_k = \nu \prod_{j < k} \theta_j$. For a uniform chain ($\theta_k = \theta$), total strategic tempo $\mathcal{T}_\Sigma$ converges to $\nu/((n+1)(1-\theta))$ as $d \rightarrow \infty$ — total strategic tempo is bounded even for arbitrarily deep chains, but the marginal contribution of each new edge decays geometrically. The deep edges are *evidence-starved*. **OR-nodes are exploration-gated**: under ε -greedy, the rate allocated to the greedy arm gets nearly all the action budget while exploratory arms get only ε/m . Pure greedy makes non-greedy arms *permanently uncorrectable*.

Per-edge persistence inherits the same weakest-link structure as per-dimension epistemic persistence (16.5): for Σ_t to persist, every edge must maintain bounded mismatch — $\forall (i, j) \in E : \nu_{ij} \cdot \iota_{ij} \cdot \eta_{\text{edge},ij} > \rho_{\Sigma,ij}/R_{\Sigma,ij}$. The aggregate $\mathcal{T}_\Sigma$ is at most the sum and at least the minimum across edges; exceeding total disturbance is *necessary but not sufficient*. Persistence is *bottleneck-limited by the weakest edge, not governed by the aggregate*. Scalar $\mathcal{T}_\Sigma$ overestimates effective strategic adaptation for the same reason scalar $\mathcal{T}$ overestimates epistemic adaptation — it averages over heterogeneous correction capacities. A practical consequence: given a fixed action budget, the topology that maximizes strategic tempo is *shallow and OR-heavy* (more edges directly observable, no attenuation); the topology that minimizes it is *deep AND-chains*. As observations accumulate per edge, $\eta_{\text{edge},ij}$ shrinks (diminishing returns), so $\mathcal{T}_\Sigma$ *declines over time even in a static environment*.

Formal Expression.

Definition (strategic-tempo).

$$\mathcal{T}_\Sigma = \sum_{(i,j) \in E} \nu_{ij} \cdot \eta_{\text{edge},ij} \cdot \iota_{ij} \quad (9.21)$$

where: - (i, j) indexes the edges of the strategy DAG (8.3) - ν_{ij} is the effective observation rate for edge (i, j) — how often the agent obtains evidence about the causal link $i \rightarrow j$ - $\eta_{\text{edge},ij}$ is the per-edge update gain (9.4) - $\iota_{ij} \in [0, 1]$ is the identifiability coefficient (9.5): the fraction of the evidence stream that genuinely identifies the edge’s causal effect

Regime contributions. The identifiability coefficient captures the regime distinction from 9.5:

Table 9.8

Regime	ι_{ij}	Contribution to $\mathcal{T}_\Sigma$	Example domain
A Intervention-rich	≈ 1	Full: $\nu_{ij} \cdot \eta_{\text{edge},ij}$	Software tests, laboratory experiments

Regime	t_{ij}	Contribution to $\mathcal{F}_\Sigma$	Example domain
B Partial intervention	$\in (0, 1)$	Reduced proportionally	Organizational actions with concurrent effects
C Observation-only	≈ 0	Near-zero: edges contribute negligibly	Passive monitoring, intelligence analysis

An agent cannot improve the parts of its strategy that it cannot test interventionally. This is the operational content of the t factor: Regime-C edges contribute essentially nothing to $\mathcal{F}_\Sigma$ regardless of how fast the agent acts or how many observations it makes. Regime-A edges yield full strategic tempo at their observation rate. The t factor is where edge-causal-validity (9.5) enters the operational machinery of strategy revision.

Parallel with epistemic tempo. The definition mirrors 3.8's $\mathcal{F} = \sum_k \nu^{(k)} \cdot \eta^{(k)*}$, replacing observation channels with strategy edges and adding the t modulation for causal identifiability. The structural parallel is exact for Regime-A edges (where $t = 1$ recovers the direct rate-times-gain form); in mixed regimes, t carries the additional content distinguishing interventional evidence from associational proxy.

Key difference: endogenous edge rates. Epistemic tempo's channel rates $\nu^{(k)}$ are largely exogenous — the environment generates observations at its own pace. Strategic tempo's edge rates ν_{ij} are *endogenous*: they depend on the agent's action policy (which edges get tested) and on upstream success (downstream edges are tested only when upstream edges fire). This endogeneity is the source of the structural differences between epistemic and strategic persistence.

Consistency verification. The definition is consistent with the four verified topologies from R:

Case B.1 (single edge $A \rightarrow G$). One edge, $\nu = \nu_{AG}$, $\eta_{\text{edge}} = 1/(n+1)$. $\mathcal{F}_\Sigma = \nu_{AG}/(n+1)$. The sector parameter $\alpha_\Sigma = 1/(n+1)$ is the per-observation correction quality; $\mathcal{F}_\Sigma = \nu \cdot \alpha_\Sigma$, matching the epistemic tempo pattern exactly.

Case B.2 (two-edge AND chain, $A \rightarrow B \rightarrow G$, B observable). Two edges. Edge 1 is tested at rate $\nu_1 = \nu$ (every execution). Edge 2 is tested only when edge 1 succeeds: $\nu_2 = \nu \cdot \theta_1$. The bottleneck edge has $\alpha_\Sigma = \min(1/(n_1+1), \theta_1/(n_2+1))$. $\mathcal{F}_\Sigma = \nu/(n_1+1) + \nu\theta_1/(n_2+1)$, consistent with depth-gated attenuation.

Case B.3 (two-edge AND chain, B unobservable). Per-edge tempo is ill-defined (the marginal point estimate is biased). Plan-level tempo is well-defined: $\mathcal{F}_{\Sigma, \text{plan}} = \nu/(n_\Phi+1)$, treating $\Phi = p_1 p_2$ as a single tracked quantity.

Case B.4 (two-arm OR node, ε -greedy). Edge 1 tested at rate $\nu_1 = \nu(1-\varepsilon)$, edge 2 at rate $\nu_2 = \nu\varepsilon$. $\mathcal{F}_\Sigma = \nu(1-\varepsilon)/(n_1+1) + \nu\varepsilon/(n_2+1)$. Action selection directly controls the rate allocation — exploration-gated, not depth-gated.

Structural decomposition. AND-chains: depth-gated (geometric attenuation). In a chain of depth d with edge success probabilities θ_k , the effective observation rate for edge k is:

Derived (Conditional on independent edges).

$$\nu_k = \nu \cdot \prod_{j < k} \theta_j \quad (9.22)$$

Each additional depth level attenuates by a factor $\theta_k < 1$. For a uniform chain ($\theta_k = \theta$, $n_k = n$ for all k):

$$\mathcal{F}_\Sigma = \frac{\nu}{n+1} \sum_{k=1}^d \theta^{k-1} = \frac{\nu}{n+1} \cdot \frac{1-\theta^d}{1-\theta} \quad (9.23)$$

This converges to $\nu/((n+1)(1-\theta))$ as $d \rightarrow \infty$ — total strategic tempo is bounded even for arbitrarily deep chains. The marginal tempo contribution of edge k decays as θ^{k-1} , falling below any fixed threshold at depth d^* (9.9). Deep AND-chains have low $\mathcal{T}_\Sigma$ at their leaves regardless of how fast the agent acts — the evidence-starvation effect identified in R.

OR-nodes: exploration-gated. At an OR-node with m alternatives under ε -exploration, the rate allocated to alternative l is:

Definition (OR-node rate allocation).

$$\nu_l = \begin{cases} \nu(1 - \varepsilon + \varepsilon/m) & l = l^* \text{ (greedy arm)} \\ \nu \cdot \varepsilon/m & l \neq l^* \text{ (exploratory arms)} \end{cases} \quad (9.24)$$

The bottleneck is the least-explored alternative. Pure greedy ($\varepsilon = 0$) gives $\nu_l = 0$ for non-greedy arms, making those edges permanently uncorrectable.

Per-edge persistence.

For Σ_t to persist, every edge must maintain bounded mismatch. The bottleneck condition is:

Derived (from persistence-condition applied per edge).

$$\forall (i, j) \in E : \quad \nu_{ij} \cdot l_{ij} \cdot \eta_{\text{edge},ij} > \frac{\rho_{\Sigma,ij}}{R_{\Sigma,ij}} \quad (9.25)$$

This is the per-edge analog of 16.5's per-dimension condition for M_t . The aggregate relationship between $\mathcal{T}_\Sigma$ and the average correction rate α_Σ is:

$$\alpha_\Sigma \leq \frac{\mathcal{T}_\Sigma}{|E|} \leq \mathcal{T}_\Sigma \quad (9.26)$$

(minimum $\leq$ average $\leq$ sum). Consequently, $\mathcal{T}_\Sigma > |E| \cdot \rho_\Sigma / R_\Sigma$ is *necessary* for persistence but not sufficient — the persistence condition is bottleneck-limited by the weakest edge, not governed by the aggregate.

Epistemic Status.

The definition itself is axiomatic — it names a quantity by analogy with 3.8. The consistency verification with the four cases from R is *derived* (conditional on Beta-Bernoulli dynamics). The AND-chain geometric attenuation is *derived* (conditional on independent edge outcomes). The OR-node exploration gating and identifiability adjustment are *hypotheses* in the general DAG case, though verified for the specific topologies above. The bottleneck-limited persistence observation is *derived* from 16.5's result applied to the strategy substate. The additive form also inherits the noise-correlation scope derived in AS: the edge-wise sum is exact when the edges' evidence-noise is cross-edge independent, and the deviation under correlated evidence (e.g., one observation informing several edges, or a shared attribution error contaminating a path) is *signed* — a common evidence source overcounts, while structurally anti-correlated evidence errors can make the sum an undercount. The per-edge topology conditions named above are independence assumptions of this kind.

Max attainable: conditional. Currently conditional because the general DAG case (mixed AND/OR topologies, correlated edges) has not been verified.

Discussion.

Connection to 16.5. The per-edge persistence condition inherits the same structure as the per-dimension epistemic result: the weak edge is the bottleneck. Scalar $\mathcal{F}_\Sigma$ overestimates effective strategic adaptation for the same reason scalar $\mathcal{T}$ overestimates epistemic adaptation — it averages over heterogeneous correction capacities. A topology-specific corollary [Discussion]: per-edge persistence (every edge calibrated) and plan-level persistence (the plan’s value estimate tracks reality) can diverge in a *stationary* environment only at OR-nodes — a starved exploratory arm drifts uncalibrated while the greedy arm carries the plan value — whereas AND-structures cannot exhibit the divergence, since every edge on the conjunctive path enters the plan value and an uncalibrated edge shows up there directly.

Three-way tradeoff. Strategic tempo competes with both epistemic tempo and exploitation for the agent’s finite action capacity. Each action that tests a strategy edge (improving $\mathcal{F}_\Sigma$) is an action not spent gathering epistemic information (improving $\mathcal{T}$) or pursuing the current best action (exploitation). The allocation is addressed in 10.2 — the extended deliberation threshold is derived, but the broader three-way framing is discussion-grade. Deliberation (internal computation) supplements $\mathcal{F}_\Sigma$ via an internal channel distinct from action-generated evidence.

Software as Regime A. In software development, tests are genuine interventions with attributable outcomes ($t_{ij} \approx 1$). This makes software agents naturally high- $\mathcal{F}_\Sigma$ — a structural advantage identified but not yet formalized in the TST domain (02-tst-core/).

Working Notes.

- **2026-07-16 — OR/AND calibration-vs-performance divergence gloss added to Discussion** (bulk-64 closeout; provenance: spikes/.integrated/spike-or-node-strategic-dynamics.md §4, whose substantive results were verified landed — this was its one unlanded discussion-grade observation).
- **2026-07-16 — noise-correlation scope made explicit (ripple from AS).** The Epistemic Status now names the cross-edge evidence-noise independence condition and its signed deviation. Nothing in the body claimed the overcount-only reading, so no claim was deleted — this closes the “same scope, unaudited” flag from the 2026-07-15 spike.
- **Mixed AND/OR DAGs.** The structural decomposition treats AND-chains and OR-nodes separately. Real strategy DAGs interleave both. How the geometric attenuation (AND) and exploration gating (OR) interact in mixed topologies is unverified. The per-edge persistence condition sidesteps this by treating each edge independently, but the *aggregate* behavior of $\mathcal{F}_\Sigma$ in mixed DAGs may exhibit interference effects (e.g., evidence starvation in an AND-chain feeding into an OR-node’s greedy arm).
- **Optimal topology.** Given a fixed action budget ν , what DAG topology maximizes $\mathcal{F}_\Sigma$? Shallow OR-heavy structures maximize tempo (more edges are directly observable, no attenuation). Deep AND-chains minimize tempo. This may yield a principled argument for preferring flat, option-rich strategies over deep sequential plans — complementing the confidence-decay argument in 8.1.
- **Dynamic complexity.** As n_{ij} grows (more observations per edge), $\eta_{\text{edge},ij}$ shrinks (diminishing returns). Strategic tempo declines over time even in a static environment — the agent’s corrections become smaller as edges converge. This is the correct qualitative behavior (converged edges need less correction) but means $\mathcal{F}_\Sigma$ is not a fixed property of the agent-environment pair.
- **Cross-reference to NeurIPS Paper 2.** The strategic-tempo aggregator is **bottleneck-not-sum** in NeurIPS 2026 Paper 2 (“A Unified Convergence Theory for Non-Stationary Reinforcement Learning”, ~/src/neurips/02-unified-convergence-rl/, §5 Lemma 4.2 / #lem-forgetting): $\mathcal{F}_\Sigma^{\text{bn,ss}} := \min_{(i,j) \in E} \nu_{ij} \cdot t_{ij} \cdot (1 - \lambda_{ij})$. Adversarial disturbance concentrates on the slowest-forgetting / weakest-attribution / slowest-revisited element — the same weakest-link pattern as #result-per-dimension-persistence at strategic-tempo level. This segment’s $\mathcal{F}_\Sigma = \sum_{(i,j)} \nu_{ij} \cdot \eta_{\text{edge},ij}$ is the throughput aggregation; the bottleneck form is the persistence-relevant aggregation. The two are compatible: throughput-sum bounds the budget of corrective work the agent can allocate; bottleneck-min bounds the rate at which the worst element survives adversarial concentration. See spikes/neurips-back-integration-2026-05-08.md §1 Paper 2 entry 8.

Incidental audit gold (lift 2026-05-31, batch A9). Cross-audit “wandering thoughts” / §14-ideation harvested from the de-novo auditors’ working dirs, deduplicated across substrates and attributed by substrate + audit. *Orthogonal* material (pedagogical framing, analogies, candidate figures, reader-confusion signals), staged for an eventual careful promotion pass, kept separate from the certified theory-fix findings. **Coverage for this segment:** 471203, 526815, 584721, 773921, 829314.

1. Candidate Brief prose / pre-prose.

- **Endogenous vs. exogenous observation rates** — “why learning a strategy is harder than learning a world model.” Independently flagged as the segment’s deepest point: epistemic observation rates are *exogenous* (the world happens to you), but strategic edge rates are *endogenous* — you only get to test a downstream edge if you succeed at the upstream ones (Claude/584721, 829314; Codex/526815). One auditor called it “the deepest architectural insight in Part II” and the source of the M_t/Σ_t persistence differences. Strong candidate Brief framing.
- **“The shape determines the physics of how the plan learns.”** A wide/OR-heavy DAG learns its paths in parallel (tempo scales with the exploration budget); a deep/AND-chain has *bounded total tempo* no matter how fast the early steps run. Candidate one-line Brief hook (Gemini/829314).

2. Candidate Discussion.

- **Bounded total tempo of deep chains** $\Rightarrow$ “waterfall is mathematically dead.” The geometric series $\sum_{k=1}^d \theta^{k-1} \rightarrow 1/(1-\theta)$ proves a deep AND-chain’s *total* learning capacity asymptotically plateaus while the marginal tempo of deep edges decays to zero — “deep plans literally cannot learn fast enough to survive a volatile environment.” Worked instance: a 20-step sequential plan starves its late steps of evidence; a startup exploring 5 verticals in parallel calibrates all five (Gemini/829314; Claude/584721, 471203). Candidate Discussion centerpiece.
- **The evolutionary reading.** Instincts are shallow OR-heavy DAGs (“if hot, pull hand away OR yell”); complex human planning is deep AND-heavy and “incredibly fragile, requiring massive cultural scaffolding (forced observability, artificially lowered ρ via institutions) to survive its own mathematical evidence starvation — we build society to lower ρ so our deep AND-chains don’t collapse before they finish” (Gemini/829314). Candidate Discussion analogy; isomorphism would need checking before promotion past discussion-grade.
- **Endogenous tempo is why “do-what-you-know” organizations fail.** Because edge rates v_{ij} are endogenous, an agent that always exploits and never tests alternatives drives v on its exploratory edges $\rightarrow 0$, so $\mathcal{F}_2$ for those edges collapses to the greedy branch only — the agent “starves its strategic learning capacity” and (combining with #der-causal-insufficiency-detection) cannot even detect its L0 insufficiency (Claude/451729). Candidate Discussion line tying the endogenous-rate property to a concrete organizational failure.
- **Five converging arguments for shallow plans.** This segment’s strategic-tempo bottleneck is the *fifth* independent pressure toward shallow strategy DAGs, joining chain-confidence-decay, evidence-starvation, identifiability-degradation, and cognitive cost. Candidate to surface as a cross-segment structural observation (Claude/584721). (*Same meta-pattern flagged under #scope-edge-update-causal-validity as the “quadruple ceiling” — this is the fifth leg; consider a single meta-segment home.*)

3. Follow-up items.

- **Dynamic complexity / gain-collapse warning.** $\mathcal{F}_2$ declines over time even in a *static* environment because $\eta_{\text{edge}} = 1/(n+1)$ shrinks as n grows — so a long-running agent eventually has insufficient strategic-correction capacity for any non-stationarity unless a ρ -spike re-inflates η . Connects to the gain-collapse / stability-induced-myopia / detection-latency-blowup pathology cluster; candidate to develop (Claude/584721; Gemini/829314). Already partly a Working-Notes item.
- **Candidate “tempo-maximizing topology” theorem (off-ramp-adjacent).** One auditor sketched that under a fixed action budget v and fixed edge count, the tempo-maximizing topology is the shallowest OR-tree: AND-chain tempo converges to $v/((n+1)(1-\theta))$ while flat-OR (full exploration) gives $v/(n+1)$, so OR-flat beats an AND-chain for depth $> 1/(1-\theta)$ (e.g. depth > 10 at $\theta = 0.9$); the Working-Notes “shallow OR-heavy maximizes tempo” line may be promotable to a derived statement (Claude/451729). *If* this firms up it is a candidate derivation (strengthen-first), not a framing bullet — flagged for the off-ramp.

- (Codex/526815 F66/F67/F68 route to the off-ramp — see lift report — as candidate findings: the headline defines $\mathcal{F}_\Sigma$ as a *sum but the persistence-relevant quantity is the bottleneck/min; “Regime-C edges contribute essentially nothing” overstates — observational evidence still improves *associational* prediction; the $\mathcal{F}_\Sigma > |E|\rho/R$ necessary bound assumes homogeneous per-edge thresholds and should sum $\sum_e \rho_e/R_e$ in the heterogeneous case.)*

4. Readers often ask / wonder.

- **How does the agent allocate a finite action budget across exploit / epistemic-explore / strategic-explore (deliberate)?** Several substrates reached forward to #disc-exploit-explore-deliberate for the three-way (or four-way, with “diagnose”) allocation rule (Claude/584721, 829314).

5. Candidate figures.

- **Throughput-sum vs. bottleneck-min, same edges.** Two simultaneous summaries over the same DAG edges: a throughput *sum* (looks healthy) and a bottleneck *minimum* (the one untested or weakly-identifiable edge that actually fails the per-edge persistence condition). Makes the sum-vs-min distinction visible at a glance (Codex/526815) — pedagogy + the F66 guardrail in one figure.

Belongs elsewhere.

- (none surfaced beyond the shared “shallow-plans meta-pattern” pointer noted under Discussion)

9.9 Formulation: Cognitive Cost of Strategy

A formulation connecting strategy structure to its *maintenance cost* — making the explicit-strategy condition from 7.6 quantitative. The minimum description length of a strategy DAG $\Sigma_t = (V, E, p, \gamma)$ decomposes into *structural* bits (encoding topology — node identities, edge connectivity, AND/OR labels) and *parameter* bits (encoding the edge credences given the topology), scaling as $O(|E| \log |V|)$ for sparse DAGs and moderate-precision credences.

The framework offers a **strategy IB objective**: Σ_t is the information-bottleneck compression of the interaction history $\mathcal{C}_t$ for guidance, parallel to M_t as the IB-compression of $\mathcal{C}_t$ for prediction. The variational form trades off compression cost $I(\mathcal{C}_t; \Sigma_t)$ against a KL divergence from the optimal-policy reference to the strategy-induced policy. The **KL direction** is *not* a stipulation — it is forced by a regret-bound derivation (T): under bounded value range and deterministic $\pi^* = \delta_{a^*}$ (AAT’s canonical scope, 6.6), the strategy-induced regret satisfies $R(Q_{\Sigma_t}) \leq V_{\max} \cdot \text{TV}(\pi^*, Q_{\Sigma_t}) \leq V_{\max} \sqrt{\frac{1}{2} D_{\text{KL}}(\pi^* \| Q_{\Sigma_t})}$ by Pinsker, with the sharper Bretagnolle-Huber identity $R \leq V_{\max}(1 - e^{-D_{\text{KL}}(\pi^* \| Q_{\Sigma_t})})$ as the primary form. The opposite-direction KL is $+\infty$ whenever Q_{Σ_t} places any mass off a^* — a vacuous bound. The direction is *forced* with π^* first (reverse-KL in variational-inference vocabulary). Within the direction-forced family, reverse-KL is *uniquely* selected by a chain-rule additivity axiom (Hobson 1969; Csiszár 1991), AAT-internally motivated as the divergence-level analog of 8.1’s log-additive layer. This is the **second instance of the four-layer additive-coordinate-forcing meta-pattern** (5.5): a uniqueness theorem operating on an AAT-internally-motivated additivity axiom forces a specific coordinate. The whole construction lifts the IB objective from discussion-grade to *robust-qualitative*.

The framework derives a **maximum useful chain depth** for AND-chains: given base observation rate ν , per-edge success probability θ , effective sample size n , and disturbance-to-capacity ratio ρ_Σ/R_Σ , there is a finite depth d^* beyond which the deepest edge cannot maintain bounded mismatch. Beyond this depth, evidence starvation makes the edge uncorrectable faster than the environment invalidates it. A central observation: **deep AND-chains suffer a triple depth penalty** that compounds. (1) *Confidence decay* (8.1): aggregate confidence $\prod p_k$ decays geometrically — *the plan is less likely to succeed*. (2) *Evidence starvation* (R): effective observation rate decays geometrically — *the plan is harder to calibrate*. (3) *Cognitive cost* (this segment): each additional depth level adds $O(\log |V|)$ bits — *the plan is more expensive to maintain*. All three are multiplicative in depth, making deep sequential strategies exponentially costly

along three independent dimensions and creating systematic pressure toward *shallow, OR-heavy strategies*. The framework also makes the maintenance cost from 7.6 concrete by decomposing it into representation cost (proportional to description length), revision cost (proportional to strategic tempo $\mathcal{T}_\Sigma$), and monitoring cost (proportional to edges with partial identifiability).

A practical consequence for language-constituted agents: a context window of W tokens imposes a hard upper bound $DL(\Sigma_t) \leq W \log_2(|\text{vocab}|)$, making the IB trade-off *non-optional* — the agent *must* compress its strategy, and d^* becomes a context-window-limited quantity (a 128K-token context may support a 500-edge DAG encoded in natural language; a 4K-token window may support only a 15-edge sketch). The Miller (2022) coevolving-automata result provides a complementary constraint from the interaction side: a complex Σ_t whose edges are only tested over a short horizon is *indistinguishable from a simpler strategy*. The interaction horizon compresses the strategy space from above, just as maintenance cost and evidence starvation compress it from below.

Formal Expression.

Strategy description length.

The minimum description length of a strategy DAG $\Sigma_t = (V, E, p, \gamma)$ (8.3) decomposes as:

Formulation (strategy-description-length).

$$DL(\Sigma_t) = DL_{\text{struct}}(G) + DL_{\text{param}}(p \mid G) \quad (9.27)$$

where: - $DL_{\text{struct}}(G)$: bits to encode the DAG topology — node identities, edge connectivity, AND/OR labels γ . Scales as $O(|E| \log |V|)$ for sparse DAGs. - $DL_{\text{param}}(p \mid G)$: bits to encode the edge credences given the topology. For Beta-distributed credences, each edge requires $O(\log n_{ij})$ bits where $n_{ij} = \alpha_{ij} + \beta_{ij}$ is the effective sample size.

The total scales as $O(|E| \log |V|)$ for moderate-precision credences, growing linearly in the number of edges and logarithmically in the number of nodes.

Strategy IB objective.

The optimal strategy complexity balances parsimony against decision-relevance. Σ_t is the IB-compression of the interaction history $\mathcal{C}_t$ for guidance, parallel to M_t as the IB-compression of $\mathcal{C}_t$ for prediction (AA for the shared IB shape across AAT's compression operations, and for the relationship between the theoretical $I(\mathcal{C}_t; \Sigma_t)$ compression cost and the operational DL-based minimization below):

Formulation (strategy-IB-objective; KL-direction strengthened by regret bound — see Epistemic Status).

Theoretical form (variational). Σ_t is a tractable variational approximation of the optimal-policy posterior $Q^*(\pi \mid M_t)$. The strategy-cost objective:

$$\Sigma_t^* = \arg \min_{\Sigma_t} \left[I(\mathcal{C}_t; \Sigma_t) + \beta_\Sigma \cdot D_{\text{KL}}(\pi^*(\cdot \mid M_t) \parallel Q_{\Sigma_t}(\pi \mid M_t)) \right] \quad (9.28)$$

where $Q_{\Sigma_t}(\pi \mid M_t)$ is the action distribution induced by the strategy DAG given the current model state, and $\pi^*(\cdot \mid M_t)$ is the optimal-policy reference. The KL direction — π^* -first — is forced by the regret-bound derivation (next paragraph); the opposite direction is vacuous under deterministic π^* .

Regret-bound derivation of KL direction. Under AAT's canonical scope, $\pi^* = \delta_{a^*}$ is deterministic (6.6). Define the strategy-induced regret against π^* as $R(Q_{\Sigma_t}) := V(a^*) - \mathbb{E}_{a \sim Q_{\Sigma_t}}[V(a)]$, where $V(a) = Q_O(M_t, a; \pi_{\text{cont}}, N_h)$ is the action-value (6.6, O_t induces V via 6.4). Under bounded value range $V_{\text{max}} := \max_a V(a) - \min_a V(a)$:

$$R(Q_{\Sigma_t}) \leq V_{\max} \cdot (1 - Q_{\Sigma_t}(a^*)) = V_{\max} \cdot \text{TV}(\pi^*, Q_{\Sigma_t}) \quad (9.29)$$

Applying Pinsker's inequality ($\text{TV}(P, Q) \leq \sqrt{\frac{1}{2} D_{\text{KL}}(P \| Q)}$) with $P = \pi^*$, $Q = Q_{\Sigma_t}$:

$$R(Q_{\Sigma_t}) \leq V_{\max} \cdot \sqrt{\frac{1}{2} D_{\text{KL}}(\pi^* \| Q_{\Sigma_t})} \quad (9.30)$$

Under deterministic π^* , $D_{\text{KL}}(\pi^* \| Q_{\Sigma_t}) = -\log Q_{\Sigma_t}(a^*)$ — finite and graded whenever $Q_{\Sigma_t}(a^*) > 0$. The opposite-direction $D_{\text{KL}}(Q_{\Sigma_t} \| \pi^*)$ equals $+\infty$ whenever Q_{Σ_t} places any mass off a^* , giving a vacuous bound. The regret-bound argument therefore **forces the KL direction** with π^* first. Within the direction-forced f-divergence family, reverse-KL is *uniquely* selected under the chain-rule additivity axiom (Hobson 1969; Csiszár 1991 Theorem 3 corollary and Theorem 5; standard functional-equation derivation per Aczél & Daróczy 1975), which is AAT-internally motivated as the divergence-level analog of additive log-confidence decay (8.1). See T §6.1 for the uniqueness theorem, §6.2 for secondary supporting characterizations (gradient-tractability, VI-alignment, MDL), and §7 for the linear-vs-square-root β_{Σ} trade-off.

The variational form is the strategy-layer analog of variational free energy minimization in active inference (Friston, FitzGerald, Rigoli, Schwartenbeck & Pezzulo 2017, “Active inference: a process theory,” *Neural Computation* 29; Da Costa, Parr, Sajid, Veselic, Neacsu & Friston 2020, “Active inference on discrete state-spaces,” *J. Math. Psych.* 99; Parr & Pezzulo 2022, *Active Inference*, MIT Press). AAT borrows the variational form as the appropriate generalization of the Shannon-MI relevance term and now derives the direction of KL from an internal regret-bound argument — without committing to AI's preferences-as-priors encoding ($C(o) = \log P_{\text{pref}}(o)$; AAT's O_t remains a value functional on trajectories, 6.4) or to expected free energy as master objective (AAT's CIY-unified objective is a related but distinct decomposition; 7.5).

Operational form. Since $I(\mathcal{C}_t; \Sigma_t)$ is not computable in closed form for general DAG encodings, the operational minimization replaces the information cost with a description-length surrogate and the KL term with a sample-based estimate (a per-edge calibration discrepancy weighted by decision-relevance — see 9.6 for the gradient form):

$$\Sigma_t^* \approx \arg \min_{\Sigma_t} \left[\text{DL}(\Sigma_t) + \beta_{\Sigma} \cdot \widehat{D}_{\text{KL}}(\pi^* \| Q_{\Sigma_t}) \right] \quad (9.31)$$

where: - $\text{DL}(\Sigma_t)$: description length (coding-cost upper bound on $I(\mathcal{C}_t; \Sigma_t)$ for the given DAG encoding scheme — see §2.2 below) - $\widehat{D}_{\text{KL}}(\pi^* \| Q_{\Sigma_t})$: sample-based estimate of the KL divergence from the optimal-policy reference to the strategy-induced policy - $\beta_{\Sigma} > 0$: trade-off parameter — cognitive cost per decision-relevant bit (the Σ_t instance of the shared β framework in AA); under the regret-bound derivation, β_{Σ} has a *local* interpretation as $V_{\max}/(2\sqrt{2D_{\text{KL}}})$ via the Pinsker form (the linear-KL form is the IB-shape instance; the square-root form is the tighter regret-scale form — see Epistemic Status and the spike for the trade-off)

The two forms agree in the limit where the DAG encoding is rate-distortion optimal and the policy posterior is sample-recoverable; the operational form is the one an agent actually runs. The theoretical form places the objective on the same variational frontier as M_t , shared intent, and composition projection, with the π^* -first KL-form relevance term resolving the Shannon-zero degeneracy under deterministic π^* and the forward-KL infinity degeneracy that the opposite direction would introduce.

When β_{Σ} is low (high maintenance cost relative to decision value), the agent prefers simple strategies. When β_{Σ} is high (strategy is cheap to maintain relative to its decision value), the agent can afford complex plans. The explicit strategy condition (7.6) is the binary threshold: β_{Σ} large enough that *any* Σ_t is worth maintaining.

Maximum useful chain depth.

From 9.8's per-edge persistence condition, an AND-chain of depth d with per-edge observation rate ν , true success probability θ per edge, and effective sample size n per edge persists only if the deepest edge satisfies:

Derived (Conditional on Beta-Bernoulli, per-edge persistence).

$$\nu \cdot \theta^{d-1} \cdot \frac{1}{n+1} > \frac{\rho_\Sigma}{R_\Sigma} \quad (9.32)$$

Solving for the maximum depth at which persistence is achievable:

$$d^* = 1 + \left\lceil \frac{\log\left(\frac{\nu}{(n+1)\rho_\Sigma/R_\Sigma}\right)}{\log(1/\theta)} \right\rceil \quad (9.33)$$

When $\nu/((n+1)\rho_\Sigma/R_\Sigma) \leq 1$, even $d = 1$ fails — the agent cannot maintain a single edge under these conditions.

Interpretation. Beyond depth d^* , evidence starvation makes edges uncorrectable faster than the environment invalidates them. The agent accumulates strategic mismatch on deep edges regardless of how fast it acts at the top of the chain.

Quantitative illustration ($\theta = 0.8, \nu = 1$):

Table 9.9

n	ρ_Σ/R_Σ	d^*
10	0.01	10
10	0.1	0
100	0.01	0
100	0.1	0

High evidence requirements (n large) and volatile environments (ρ_Σ/R_Σ large) severely limit useful chain depth.

Triple depth penalty. Deep AND-chains suffer three independent penalties that compound:

1. **Confidence decay** (8.1): aggregate confidence $\prod p_k$ decays geometrically with depth. The plan is *less likely to succeed*.
2. **Evidence starvation** (R): effective observation rate $\nu_k = \nu \cdot \prod_{j < k} \theta_j$ decays geometrically. The plan is *harder to calibrate*.
3. **Cognitive cost** (this segment): each additional depth level adds $O(\log|V|)$ bits to description length. The plan is *more expensive to maintain*.

All three are multiplicative in depth, making deep sequential strategies exponentially costly along three independent dimensions.

Enriched explicit strategy condition.

The maintenance cost C_{maintain} from 7.6 decomposes as:

Formulation (enriched-strategy-condition).

$$C_{\text{maintain}} = C_{\text{represent}} + C_{\text{revise}} + C_{\text{monitor}} \quad (9.34)$$

where: - $C_{\text{represent}} \propto \text{DL}(\Sigma_t)$: cognitive cost of holding the strategy in working memory (proportional to description length) - $C_{\text{revise}} \propto \sum_{(i,j)} \nu_{ij} \cdot c_{\text{update}}$: cost of processing edge updates (proportional to strategic tempo $\mathcal{T}_\Sigma$ times per-update cost) - $C_{\text{monitor}} \propto |\{(i,j) : \iota_{ij} < 1\}|$: cost of monitoring edges with partial identifiability (the agent must do extra causal reasoning for non-trivial edges)

This decomposition makes the 7.6's maintenance term concrete: each component maps to a quantity defined elsewhere in the theory.

Complexity compression operations.

The IB objective suggests three compression operations, corresponding to structural changes from 9.7:

Discussion (complexity-compression).

1. **Edge pruning** (operation 3 in 9.7): remove edges with $\eta_{\text{edge},ij} \cdot I_{\text{edge},ij} < c_{\text{bit}}$, where $I_{\text{edge},ij}$ is the decision-relevance of that edge and c_{bit} is the per-bit maintenance cost. Edges that contribute less decision value than their representational cost are candidates for removal.
2. **Node merging** (reducing $|V|$): collapse intermediate nodes that serve no decision-distinguishing function. This reduces $\text{DL}_{\text{struct}}$ by a factor proportional to the reduction in $|V|$.
3. **Depth truncation** at d^* : prune all edges beyond the maximum useful depth. This is not optimization but necessity — edges beyond d^* cannot maintain bounded mismatch.

Epistemic Status.

The description length formulation is a *formulation* — it applies standard MDL to the strategy DAG, which is a representational choice not a derived necessity. The IB objective is *formulation (strengthened by regret bound)* in its variational form (above): the specific KL *direction* (π^* -first, i.e., reverse-KL in the variational-inference vocabulary) is *derived* as an upper regret bound via Pinsker's inequality under bounded value range and deterministic π^* (see Regret-bound derivation paragraph above; full derivation in appendix T). The choice of reverse-KL *within* the direction-forced family upgrades from canonical-formulation to **derived (conditional on chain-rule additivity axiom)** — the chain-rule axiom (Hobson 1969; Csiszár 1991 Theorem 3 corollary and Theorem 5; standard functional-equation derivation per Aczél & Daróczy 1975) picks reverse-KL uniquely among f-divergences, and the axiom is AAT-internally motivated as the divergence-level analog of 8.1 (see T §6.1). Secondary properties (gradient-tractability, variational-inference alignment with Friston et al. 2017, Da Costa et al. 2020, Parr & Pezzulo 2022; MDL coding; compatibility with Amari & Nagaoka 2000 Fisher geometry) are convergent evidence rather than independent uniqueness grounds — in particular, Fisher-metric-at-second-order is *not* distinguishing within f-divergences (Eguchi 1983, *Ann. Statist.* 11(3):793–803). The regret-bound derivation closes the direction ambiguity that the earlier V-medium move (commit a14682e) left open: the initial V-medium form used $D_{\text{KL}}(Q_{\Sigma_t} \parallel \pi^*)$ (forward-KL), which is $+\infty$ under deterministic π^* whenever Q_{Σ_t} has any off-optimum mass — a different-valued but structurally identical degeneracy to the Shannon-MI zero it replaced. The π^* -first reverse-KL direction escapes both degeneracies by construction.

The variational form replaces an earlier Shannon-MI form $-\beta_\Sigma \cdot I(\Sigma_t; \pi^* \mid M_t)$ which had a Shannon-zero degeneracy: when π^* is a deterministic function of M_t (the standard scope), Shannon mutual information to a constant vanishes

identically, collapsing the objective to $\arg \min \text{DL}(\Sigma_t)$. The π^* -first KL form does not have this degeneracy and is graded whenever Q_{Σ_t} places any mass on a^* . The maximum useful depth d^* is *derived* conditional on Beta-Bernoulli dynamics and the per-edge persistence condition from 9.8. The triple depth penalty is an *observation* combining results from three independent segments. The enriched maintenance decomposition is *formulation*. The compression operations are *discussion-grade*.

On β_Σ interpretation. Under the Pinsker-reverse-KL bound $R(Q_{\Sigma_t}) \leq V_{\max} \sqrt{\frac{1}{2} D_{\text{KL}}(\pi^* \| Q_{\Sigma_t})}$, a square-root-in-KL trade-off would naturalize β_Σ globally as $\beta_\Sigma \propto V_{\max}$. The segment retains the linear-in-KL form (to preserve the rate-distortion-Lagrangian IB shape shared with AA); under the linear form, β_Σ has only a *local* regret-bound interpretation ($\partial R / \partial D_{\text{KL}}$ at the operating point). This is a trade-off between IB-shape alignment and regret-scale naturalization. Under AAT’s canonical deterministic- π^* scope, the sharper Bretagnolle-Huber identity $R \leq V_{\max}(1 - e^{-D_{\text{KL}}(\pi^* \| Q_{\Sigma_t})})$ holds as the primary form (T §4); Pinsker is retained here for IB-shape alignment and as the loose general form correct for stochastic- π^* extensions.

Assumption explicitly stated: bounded value range. The regret-bound derivation requires $V_{\max} := \max_a V(a) - \min_a V(a) < \infty$ over $\mathcal{A}$ at fixed M_t . This is mild but not automatic — 6.4 specifies $V_{O_t} : \text{trajectories} \rightarrow \mathbb{R}$, and bounded range at fixed state is an additional assumption stated here.

Max attainable: *robust-qualitative* for the IB objective with the direction-forced derivation; conditional for the specific functional form (linear vs. square-root in KL). The DL formulation is standard; the depth bound could reach exact status for specific edge models. The regret-bound derivation does not extend to stochastic π^* (outside AAT canonical scope); see 6.6 continuation conventions for scope.

Discussion.

Connection to 7.6. The enriched maintenance decomposition gives the cost inequality quantitative content: an agent can now *compute* whether a proposed strategy is worth maintaining by estimating its description length, revision cost, and monitoring burden, and comparing against the exploration/repair cost of operating without it.

LLM context windows as DL constraint. For language-constituted agents (03-llm-core/), the strategy must fit in the context window. A context window of W tokens imposes $\text{DL}(\Sigma_t) \leq W \cdot \log_2(|\text{vocab}|)$ as a hard constraint. This makes the IB trade-off non-optional: the agent *must* compress its strategy, and the depth bound d^* becomes a context-window-limited quantity. A 128K-token context window may support a 500-edge DAG encoded in natural language; a 4K-token window may support only a 15-edge sketch.

Computational compression from interaction horizon (Miller 2022). The maximum useful depth d^* derived above constrains complexity from the *maintenance* side (evidence starvation). Miller (2022, *Ex Machina*, Table 12.2) provides a complementary constraint from the *interaction* side: the number of interaction rounds compresses the space of behaviorally distinguishable strategies regardless of the agent’s internal complexity. For Moore machines with binary actions:

Table 9.10

<i>Agent states</i>	<i>Unique computations</i>	<i>After 1 round</i>	<i>After 2 rounds</i>	<i>After 4 rounds</i>
1	2	2	2	2
2	26	2	8	26
3	1,054	2	8	690
4	57,068	2	8	5,936

The pattern: effective complexity = min(agent complexity, interaction-horizon complexity). With only one round, even a four-state machine (57,068 unique computations) reduces to two distinguishable behaviors — equivalent to a one-state machine. With two rounds, all machines with two or more states reduce to eight distinguishable behaviors. The interaction horizon compresses the strategy space from above, just as the maintenance cost and evidence starvation compress it from below. For AAT’s strategy DAG, the analog: a complex Σ_t whose edges are only tested over a short horizon gains nothing from its depth — the untested structure is indistinguishable from a simpler strategy. This reinforces the d^* bound and provides empirical grounding beyond the Beta-Bernoulli derivation.

Strategy simplification pressure. The triple depth penalty creates systematic pressure toward shallow, wide (OR-heavy) strategies over deep, sequential (AND-heavy) ones. This aligns with the structural pressure identified in 8.1’s Discussion, now grounded in three independent mechanisms rather than one.

Working Notes.

- **Mixed topologies.** The depth bound d^* assumes uniform AND-chains. Mixed AND/OR DAGs have heterogeneous depth penalties: OR-nodes reset the evidence-starvation clock (each alternative is tested independently), while AND-nodes compound it. The effective depth for computing d^* may be the longest AND-chain in the DAG, not the total graph depth.
- **Optimal topology.** Given a fixed DL budget and action rate ν , what DAG topology maximizes decision-relevant information $I(\Sigma_t; \pi^* | M_t)$? This is a combinatorial optimization over graph structures — likely NP-hard in general but potentially tractable for specific graph families (trees, bounded treewidth).
- **Dynamic complexity.** As edges converge (high n_{ij}), their per-edge $\eta_{\text{edge},ij}$ shrinks, but their description length DL_{param} grows (more bits to encode the precise credence). The IB objective would favor *dropping* converged edges (they contribute little decision-relevant information since the agent already acts correctly on them), replacing them with a default “high confidence” summary. This is a principled version of “stop tracking what you already know.”
- **Stochastic $\mathcal{F}_\Sigma$.** If strategic disturbance follows Model S (stochastic) rather than Model D (deterministic), the steady-state mismatch scales as $\rho_\Sigma / \sqrt{\mathcal{F}_\Sigma}$ rather than $\rho_\Sigma / \mathcal{F}_\Sigma$. The depth bound d^* would change accordingly. Not yet derived.

Incidental audit gold (lift 2026-05-31, batch A9). Cross-audit “wandering thoughts” / §14-ideation harvested from the de-novo auditors’ working dirs, deduplicated across substrates and attributed by substrate + audit. *Orthogonal* material (pedagogical framing, analogies, candidate figures, reader-confusion signals), staged for an eventual careful promotion pass, kept separate from the certified theory-fix findings. **Coverage for this segment:** 526815, 584721, 773921, 829314, 849201. (Multiple substrates rated this segment top-decile.)

1. Candidate Brief prose / pre-prose.

- **“You cannot out-plan the entropy of the environment.”** The plain-language stake of the d^* max-useful-depth result: beyond d^* , “the agent accumulates strategic mismatch on deep edges regardless of how fast it acts at the top of the chain” — a “devastating mathematical critique of 5-year corporate strategic plans” (Gemini/829314). Strong candidate Brief hook for the depth bound.
- **“Don’t plan further than you can reliably update your plan.”** The one-line gloss for bounded-rationality-in-planning (Claude/849201). And: over-planning is “not a psychological flaw but a mathematical error that burns adaptive reserve maintaining uncorrectable beliefs” (Gemini/773921).

2. Candidate Discussion.

- **The regret-bound derivation forcing reverse-KL — frame it as the headline strengthening.** Several substrates singled this out as a recent, clean structural win: it converts an arbitrary-looking KL *direction* choice into a decision-theoretic *forcing* argument and resolves the Shannon-zero degeneracy (MI to a deterministic π^* vanishes) and the forward-KL $+\infty$ degeneracy at once (Claude/584721, 849201; Gemini/773921, 829314). Candidate to foreground in Discussion as “direction is derived, not stipulated.”

- **LLM context-window as a literal DL budget.** The $W \log_2 |\text{vocab}|$ bound on $DL(\Sigma_t)$ is “the strongest connection between AAT and modern GenAI engineering in the framework so far” (Gemini/773921). Vivid reader-facing consequences: a chain-of-thought plan exceeding the window induces “sudden-onset dementia — the agent forgets why it is doing what it is doing” (Gemini/829314); and a *counterintuitive prediction* worth stating — small-context agents naturally adopt shallow/wide (OR-heavy) strategies, while large-context agents “might fall into the trap of building brittle deep AND-heavy plans that fail due to volatility before they finish” (Gemini/829314; Claude/526815, 584721). Candidate Discussion bridge to 03-llm-core/.
- **Six converging arguments for shallow plans (the meta-pattern).** This segment consolidates the *triple* depth penalty (confidence-decay + evidence-starvation + cognitive-cost); adding the strategic-tempo bottleneck (#def-strategic-tempo), identifiability-degradation (#scope-edge-update-causal-validity), and Miller’s interaction-horizon compression from *above* yields six independent mechanisms compounding to the same conclusion. Flagged as a candidate organizing-principle observation worth elevating to a meta-segment (Claude/584721). (*Same meta-pattern surfaced under #scope-edge-update-causal-validity and #def-strategic-tempo — one home, not three.*)

3. Follow-up items.

- **“Stop tracking what you already know” — compression-by-convergence.** The IB objective favors *dropping* converged (high- n) edges, whose precise credences cost more bits while contributing little decision-relevant information, replacing them with a “high confidence” default — a principled (not heuristic) continual-learning operation (Claude/584721). Already a Working-Notes bullet; flagged as a candidate to develop.
- **Resolved during the audit window:** Codex/526815 (F69) flagged the d^* illustration table as arithmetically wrong (claimed $d^* = 5$ for the $n = 100, \rho/R = 0.01$ case where the formula gives 0); the current table reads $d^* = 0$ for that row (verified 2026-05-31) — the correction has landed. Similarly the Bretagnolle-Huber identity is now the *primary* regret form in canon with Pinsker shown as the intermediate, largely discharging 584721’s F-D2 “Pinsker-only” integration-debt note. (*The remaining Codex F70 — C_revise scaling as raw v_{ij} vs strategic-tempo — and F71 — undeclared warrant-bearing deps — route to the off-ramp as candidate findings; see lift report.*)

4. Readers often ask / wonder.

- **How does an agent actually perform “node merging” on a causal DAG?** Is it abstracting several concrete actions into one macro-action (the Options framework / hierarchical RL)? The compression operations (edge-pruning, node-merging, depth-truncation) invite this question (Claude/829314, 849201).
- **Editorial signal:** two substrates read the Active-Inference-differentiation paragraph and the Pinsker-vs-Bretagnolle-Huber detail as belonging in the appendix (#deriv-strategy-cost-regret-bound), leaving the operational objective + Pinsker intuition in-body (Claude/829314). Preserved as a placement signal for the eventual pedagogy pass.

5. Candidate figures.

- **Sandwich-bound diagram.** Strategy complexity bounded *from below* by maintenance cost / evidence starvation (d^*) and *from above* by Miller’s interaction-horizon compression (effective complexity = $\min(\text{agent complexity}, \text{horizon complexity})$) — “even if you can afford a deep DAG, a short interaction horizon means the depth doesn’t pay off” (Claude/584721). A two-sided-bound figure communicates this faster than the prose.
- **Triple-depth-penalty panel.** The three compounding mechanisms (decay / starvation / cost) as a single multiplicative-in-depth visual; “should be visually highlighted as a major structural result” (Gemini/829314).

Belongs elsewhere.

- **Six-mechanism shallow-plan convergence $\Rightarrow$ a meta-segment.** As above, this cross-segment structural observation (spanning der-chain-confidence-decay, der-observability-dominance, scope-edge-update-causal-validity, def-strategic-tempo, this segment, and Miller’s horizon bound) is a candidate organizing principle for the Meta-Architecture cluster rather than living implicitly across six segments (Claude/584721).

- **Context-window-driven strategy-shape prediction** → 03-11m-core/. The prediction that effective context budget shapes whether a logogenic agent adopts shallow-wide vs deep-brittle strategies is reach into the logogenic-agent material, not this segment (Gemini/829314; Claude/526815).

9.10 Proposed Schema: Proposed-schema: Strategy Persistence Schema

The strategic-layer instantiation of the sector-persistence template (G). The template proves bounded state for any system with a state variable, a correction function satisfying the sector condition, and bounded disturbance — it is domain-agnostic across its preconditions (T1)–(T3). The strategic instance asks whether the same machinery transfers to strategy update dynamics, and the answer is: *yes, but conditionally*, with the conditions tight enough that the schema’s reachable persistence region has a λ -independent structural ceiling. The structural identity with the epistemic persistence condition from Part I is not an analogy; it is a *mathematical identity at the sector-framework level* — and the strategic analog of “the environment changes faster than the agent can learn” is “the world invalidates plans faster than the agent can revise them,” producing the same catastrophic outcome.

This is status: conditional, conditioned on four named premises: (i) Beta-Bernoulli edge dynamics, (ii) exponential forgetting with discount $\lambda \in (0, 1)$, (iii) the template preconditions (T1)–(T3), and (iv) a mismatch state chosen from the per-edge credence error δ_c or the plan-level surrogate $\delta_s = \hat{R}_\Sigma - \Phi$ (credit-assignment-free; persistence proved in R Prop B.5). Within these conditions, the exact steady-state sector parameter is $\alpha_\Sigma^{ss} = (1 - \lambda)/(2 - \lambda)$, and the persistence prerequisite is $(1 - \lambda)/(2 - \lambda) > \rho_\Sigma/R_\Sigma$. The linear form $1 - \lambda$ is the slow-forgetting asymptote ($\lambda \rightarrow 1$), valid only in that limit; outside the asymptote it overestimates the steady-state by up to a factor of 2.

The strategic-side asymmetry from the epistemic case — and what forces the conditional tier — is **experience discounting as structural prerequisite**: the strategic sector parameter has the form $\alpha_\Sigma = 1/(n+1)$ in raw Beta-Bernoulli, which decays monotonically with experience and drives any agent across the persistence threshold asymptotically regardless of disturbance rate. Exponential forgetting stabilizes the effective sample size at $1/(1 - \lambda)$, giving the exact $(1 - \lambda)/(2 - \lambda)$ form. The forgetting prerequisite is therefore a *trajectory-guarantee requirement*, not a tunable hyperparameter — an agent without forgetting has no long-run strategic persistence. The schema’s reachable persistence region under any λ is bounded above by a hard structural ceiling at $\rho_\Sigma = R_\Sigma/2$, derived self-contained in S: when strategic disturbance reaches half the strategic reserve, no λ satisfies the prerequisite, and the schema’s trajectory guarantee fails regardless of forgetting design. Strategic mismatch under the credit-assigned form $\delta_{\text{strategic}}$ (9.1) remains open (see Discussion §1 and 9.6).

Formal Expression.

If strategic update dynamics satisfy the template preconditions (T1)–(T3) of G for $\xi = \delta_\Sigma$ (a strategic mismatch state), together with:

Proposed Schema (strategy-persistence-schema, from sector-persistence-template).

- **(SA1)** Zero correction at zero strategic mismatch (the template’s (T1)): when the mismatch state is zero, no revision occurs
- **(SA2’)** Local sector condition on strategic correction (the template’s (T2) for $\xi = \delta_\Sigma$): the correction function points inward with baseline efficiency α_Σ within a strategic reserve R_Σ
- **(SA3)** Sufficient exploration (OR-nodes only): the action selection policy allocates correction capacity to all OR alternatives at a rate exceeding the strategic disturbance-to-reserve ratio
- Bounded strategic disturbance at rate ρ_Σ (the template’s (T3)): the rate at which the environment invalidates causal links is bounded

Then Σ_t persists iff:

$$\alpha_\Sigma > \frac{\rho_\Sigma}{R_\Sigma} \quad (9.35)$$

directly by the template's Model D result. Here α_Σ is the strategic correction rate, ρ_Σ is the strategic disturbance rate, and R_Σ is the strategic reserve (tolerance for strategic mismatch before performance degrades catastrophically). The Model S instantiation replaces ρ_Σ with σ_Σ and gives $\alpha_\Sigma > n\sigma_\Sigma^2/(2R_\Sigma^2)$ under the same template.

Forgetting as Prerequisite.

The schema form above is an **instantaneous persistence check at the current experience level**, not a trajectory guarantee. For Beta-Bernoulli edge updates (the canonical verified case; see R Props B.1–B.6), the sector parameter has the form:

Formulation (forgetting-prerequisite).

$$\alpha_\Sigma = \frac{1}{n+1} \quad (9.36)$$

where n is the edge's accumulated experience (pseudo-count). Without a forgetting mechanism, n grows monotonically with each observation, so $\alpha_\Sigma \rightarrow 0$ for every edge asymptotically. For any fixed (ρ_Σ, R_Σ) with $\rho_\Sigma > 0$, every agent eventually violates the threshold. The structural identity with 4.5 — where α can be stationary — holds for the strategic case only under an explicit forgetting mechanism.

Exponential forgetting. Replace the raw Beta-Bernoulli update with a discounted update: at each step, shrink the pseudo-counts by a factor $\lambda \in (0, 1)$:

$$\alpha_k \mapsto \lambda \alpha_k + y_k, \quad \beta_k \mapsto \lambda \beta_k + (1 - y_k) \quad (9.37)$$

The effective sample size stabilizes at $n_{\text{eff}} = 1/(1 - \lambda)$, and substituting into Prop B.1's $\alpha_\Sigma = 1/(n+1)$ gives the exact steady-state sector parameter:

$$\alpha_\Sigma^{\text{ss}} = \frac{1}{n_{\text{eff}} + 1} = \frac{1 - \lambda}{2 - \lambda} \quad (9.38)$$

For slow forgetting ($\lambda \rightarrow 1$, the regime where the prerequisite is most likely to bind), the leading-order expansion gives the simpler form $\alpha_\Sigma^{\text{ss}} \approx 1 - \lambda$ — which agrees with the forgetting rate itself. Outside the high- λ regime the approximation deteriorates: at $\lambda = 0.5$, the exact form gives $\alpha_\Sigma^{\text{ss}} = 1/3$ while the linear approximation gives $1/2$ ($\approx 50\%$ overestimate); at $\lambda = 0.9$, exact $\alpha_\Sigma^{\text{ss}} = 1/11$ versus approximation $1/10$ ($\approx 10\%$ overestimate).

The forgetting prerequisite. Combining with the schema's persistence form $\alpha_\Sigma > \rho_\Sigma/R_\Sigma$:

$$\frac{1 - \lambda}{2 - \lambda} > \frac{\rho_\Sigma}{R_\Sigma} \iff \lambda < \frac{R_\Sigma - 2\rho_\Sigma}{R_\Sigma - \rho_\Sigma} \quad (9.39)$$

(valid when $\rho_\Sigma < R_\Sigma/2$; for $\rho_\Sigma \geq R_\Sigma/2$ no λ satisfies the prerequisite and the schema's trajectory guarantee fails for any forgetting rate). The hard ceiling at $\rho_\Sigma = R_\Sigma/2$ and the algebraic content of the steady-state form are derived self-contained in S — a λ -independent structural cap on the schema's reachable persistence region under any exponential-forgetting design.

This is a **prerequisite of the schema's trajectory guarantee, not a tunable heuristic**. An agent without forgetting has no long-run strategic persistence regardless of its initial α_Σ . The steady-state sector parameter must exceed the disturbance-to-reserve ratio, or the instantaneous persistence check — no matter how comfortably it holds at any given time — eventually fails as experience accumulates.

In the slow-forgetting regime the threshold simplifies to the linear-form analog of 4.5:

$$(1 - \lambda) > \frac{\rho_\Sigma}{R_\Sigma} \quad (\text{slow-forgetting limit, } \lambda \rightarrow 1) \quad (9.40)$$

playing the role that $\mathcal{T} > \rho/\|\delta_{\text{critical}}\|$ plays for the epistemic case. The forgetting rate $(1 - \lambda)$ is the strategic analog of adaptive tempo: faster forgetting means faster tracking but noisier estimates; slower forgetting means stable estimates but slower tracking. The optimal λ balances bias and variance for the specific ρ_Σ the environment presents. Outside the slow-forgetting regime, the exact form $(1 - \lambda)/(2 - \lambda) > \rho_\Sigma/R_\Sigma$ is the operating threshold and the simpler linear form is unsafe (it permits λ values that violate the actual prerequisite).

Which mismatch state? The schema applies to any mismatch state for which conditions (SA1)-(SA3) can be verified. Two candidates exist:

- **Strategy-plan-confidence error** $\delta_s = \hat{R}_\Sigma - \Phi$: the scalar difference between the agent's strategy-plan-confidence score and the independence-model plan value at true edge parameters. This is the mismatch for which persistence IS proved (Props B.1-B.5 in R). It is computable from status propagation without credit assignment. **Scope:** δ_s operates at L0 of the Correlation Hierarchy (8.3) — it tracks calibration within the independence model. For L1 (augmented DAG), the same persistence result applies to the augmented graph's $\hat{R}_\Sigma$. The gap between L0's Φ and actual plan success under correlated failure is a model-class limitation, not an estimation error.
- **Strategic-calibration residual** $\delta_{\text{strategic}}$: the per-edge value-increment residual aggregation defined in 9.1. This is the mismatch the orient cascade (10.1) uses for edge-level revision. Persistence of $\delta_{\text{strategic}}$ remains **open** and requires the credit-assignment machinery in 9.6.

The verified instances below all use per-edge credence error $\delta_c = (\hat{p}_k - \theta_k)$ or the plan-level surrogate δ_s . They do not verify the schema for $\delta_{\text{strategic}}$ directly.

Epistemic Status.

Conditional, conditioned on (i) Beta-Bernoulli edge dynamics, (ii) exponential forgetting with $\lambda \in (0, 1)$, (iii) the sector-persistence template preconditions (T1)-(T3) of G, and (iv) the mismatch state chosen from $\{\delta_c, \delta_s\}$ — the per-edge credence error or the plan-level surrogate (per the dependency on R Props B.1-B.6). Within these conditions the exact threshold $(1 - \lambda)/(2 - \lambda) > \rho_\Sigma/R_\Sigma$ and the hard ceiling at $\rho_\Sigma \geq R_\Sigma/2$ are algebraically exact — derived self-contained in S. Strategic mismatch under $\delta_{\text{strategic}}$ (the credit-assigned form) remains open; see Discussion §1 and 9.6.

What was missing in the original sketch was instantiation — showing that specific strategic update dynamics satisfy the template's preconditions. Four cases have now been verified (full derivations in R):

1. **Single edge, Beta-Bernoulli** ($A \rightarrow G$): Sector condition satisfied globally with $\alpha_\Sigma = 1/(n + 1)$. The bound is tight (expected correction is exactly linear). (A1) satisfied. Persistence condition: $1/(n + 1) > \rho_\Sigma/R_\Sigma$.
2. **Two-edge chain, observable intermediate** ($A \rightarrow B \rightarrow G$, B observable): Sector condition satisfied globally with $\alpha_\Sigma = \min(1/(n_1 + 1), \theta_1/(n_2 + 1))$ — a weakest-link result. Correction function is diagonal (no cross-edge coupling). (A1) satisfied. The θ_1 factor in edge 2's rate is the evidence-starvation effect.

3. **Two-edge chain, unobservable intermediate** ($A \rightarrow B \rightarrow G, B$ not observable): Per-edge sector condition **fails** — the marginal Bayesian update violates (A1) with bias $O(1/n)$. But plan-level tracking (treating $\hat{\Phi} = p_1 p_2$ as a single Beta) recovers the sector condition with $\alpha_{\Sigma, \text{plan}} = 1/(n_{\Phi} + 1)$, at the cost of per-edge diagnostic resolution.
4. **Two-arm OR-node, ε -greedy** ($A_1, A_2 \rightarrow G, G$ is OR): Sector condition satisfied with $\alpha_{\Sigma} = \min((1 - \varepsilon)/(n_1 + 1), \varepsilon/(n_2 + 1))$ — an **exploration-gated** weakest-link, not depth-gated as in AND chains. (SA3) required: minimum exploration rate $\varepsilon > \rho_{\Sigma}(n_{\max} + 1)/R_{\Sigma}$. Pure greedy ($\varepsilon = 0$) **fails** the sector condition. With optimal equal-rate allocation, $\alpha_{\Sigma} = 1/(n_1 + n_2 + 2)$ — the correction capacity is split across alternatives.

The schema is no longer purely hypothetical. The sector parameter for strategic dynamics is the edge update gain η_{edge} — the same quantity that governs epistemic persistence. The structural parallel between epistemic and strategic persistence is not an analogy but a mathematical identity at the sector-framework level.

Discussion.

What's needed to promote this from schema to result.

1. **Strategic mismatch state:** partially resolved. Prop B.5 in R shows the sector condition transfers from per-edge credence error to **strategy-plan-confidence error** $\delta_s = \hat{R}_{\Sigma} - \Phi$ — the scalar difference between the agent's strategy-plan-confidence score and the independence-model plan value at true edge parameters (note: Φ is NOT actual plan success probability under correlated failure — see 8.3 edge-independence caveat). For linear correction (Beta-Bernoulli), the transfer is exact ($\alpha_s = \alpha_c$); for nonlinear correction, $\alpha_s \geq \alpha_c/\kappa(\mathbf{J})^2$. **However**, δ_s is distinct from the **strategic-calibration residual** $\delta_{\text{strategic}}$ defined in 9.1, which is an L^2 aggregation of per-edge value-increment residuals requiring credit assignment to compute. Persistence of δ_s (plan-level tracking) is proved; persistence of $\delta_{\text{strategic}}$ (per-edge diagnostics) remains open and requires the credit-assignment machinery in 9.6.
2. **Strategic correction function:** needs to satisfy the sector condition. **Resolved** for Beta-Bernoulli edges. Props B.1-B.4 in R verify the sector condition for four topologies (single edge, two-edge AND observable/unobservable, two-arm OR).
3. **Strategic disturbance:** The rate at which the environment invalidates causal links in Σ_t . **Still open** as a formalized quantity — currently a domain parameter (ρ_{Σ}), analogous to how ρ for epistemic disturbance is a domain parameter in 4.5.
4. **Sector condition verification:** the critical mathematical step. **Resolved** for four topologies. See R.
5. **Credit assignment / signal function:** needed for edge updates. **Characterized at the theory level.** 9.6 shows persistence does not require credit assignment (Prop B.5), establishes directional fidelity as the minimal requirement, and provides a gradient-based default signal function. The specific update algorithm is domain engineering, not theory — the same way the gain *estimator* is domain engineering while the gain *principle* ($\eta^* = U_M/(U_M + U_o)$) is theory. Caveat: the default gradient signal inherits $\hat{R}_{\Sigma}$'s overestimation bias under correlated failures (8.3, 9.6).
6. **Time-varying α_{Σ} :** this is where the strategic case genuinely differs from the epistemic one. For Beta-Bernoulli edges, $\alpha_{\Sigma} = 1/(n + 1)$ decays monotonically with experience, so the sector-persistence template's constant- α precondition cannot hold asymptotically under any raw Bayesian update. The resolution is the **forgetting prerequisite** promoted into the Formal Expression above: exponential forgetting at rate λ stabilizes α_{Σ} at the exact value $(1 - \lambda)/(2 - \lambda)$, with the simpler linear form $1 - \lambda$ as the slow-forgetting asymptote. The exact threshold $(1 - \lambda)/(2 - \lambda) > \rho_{\Sigma}/R_{\Sigma}$ converts the schema's instantaneous check into a trajectory guarantee. This is structural, not heuristic — without forgetting, the schema's form holds only until the agent accumulates enough experience to cross below threshold; with insufficient forgetting (linear-form approximation that treats ρ_{Σ} as small relative to R_{Σ} when in fact $\rho_{\Sigma} \geq R_{\Sigma}/2$), the prerequisite is unsatisfiable and the schema's trajectory guarantee fails.

The structural parallel is genuine, but conditional on forgetting. This schema extends *structural persistence* (see Persistence in LEXICON.md) from the epistemic substate to the strategy substate — asking whether the strategy correction machinery can outpace the rate at which the environment invalidates the agent’s causal theory. It inherits the same limitation: structural persistence of Σ_t does not address operational persistence (how close $\|\delta_{\text{strategic}}\|$ is to R_Σ) or continuity persistence (whether the agent’s strategic identity coheres through time). The persistence condition for M_t (4.5) says: adaptive tempo must exceed the ratio of disturbance to critical mismatch. If the same mathematics applies to Σ_t , then strategy persistence requires strategic tempo to exceed the ratio of strategic disturbance to critical strategic mismatch. The strategic analog of “the environment changes faster than the agent can learn” is “the world invalidates plans faster than the agent can revise them.” Both lead to the same catastrophic outcome: the system cannot maintain bounded mismatch and begins to degrade.

What this would buy the theory. If promoted to a result, strategy persistence would: - Provide a formal criterion for “when does a strategy remain viable?” - Connect strategic failure modes to the same mathematical framework as epistemic failure modes - Enable quantitative comparison: is the bottleneck epistemic persistence (model can’t keep up with reality changes) or strategic persistence (plans can’t keep up with requirement changes)? - Ground the organizational intuition that plans need to be revised faster than the situation changes

Findings.

The Forgetting Prerequisite for Strategic Persistence. Brief: With infinite memory, standard Bayesian updating guarantees eventual strategic failure. The sector parameter for Beta-Bernoulli edge updates is $\alpha_\Sigma = 1/(n + 1)$, which decays monotonically as experience n accumulates. For any positive disturbance rate, the persistence threshold $\alpha_\Sigma > \rho_\Sigma/R_\Sigma$ is eventually violated regardless of the agent’s initial calibration. Exponential forgetting with discount factor λ stabilizes the effective sample size at $1/(1 - \lambda)$, giving exact steady-state $\alpha_\Sigma^{\text{ss}} = (1 - \lambda)/(2 - \lambda)$ — with the simpler linear form $1 - \lambda$ as the slow-forgetting asymptote ($\lambda \rightarrow 1$). The forgetting prerequisite — $(1 - \lambda)/(2 - \lambda) > \rho_\Sigma/R_\Sigma$ — converts the schema’s instantaneous persistence check into a trajectory guarantee. Forgetting is therefore a structural prerequisite, not a tunable heuristic: the rate at which the agent discounts old evidence must satisfy the threshold. The schema’s reachable persistence region under any λ is bounded above by a hard structural ceiling at $\rho_\Sigma = R_\Sigma/2$, derived self-contained in #deriv-strategic-persistence-hard-ceiling.

Impact: Translates an organizational platitude (“stay adaptive”) into a quantitative survival inequality with explicit failure mode. Identifies a structural calcification process common to all long-running adaptive systems whose update mechanism accumulates evidence: institutional rigidity, RL value-function staleness, scientific-paradigm lock-in, and the loss-of-edge phenomenon in incumbent firms all instantiate the same dynamic. The threshold is sharp — no amount of prior learning protects against an environment that invalidates plans faster than the agent forgets. For long-lived AI agents this mandates explicit memory-pruning mechanisms; the design choice is not whether to forget but at what rate, and how to keep the strategic disturbance well below half the reserve (above which the hard ceiling closes the design space — see #deriv-strategic-persistence-hard-ceiling). The structural identity with #result-persistence-condition (where α can be stationary) is recovered for the strategic case *only* under explicit forgetting, which is itself a non-obvious finding about the asymmetry between epistemic and strategic dynamics.

Novelty Claim: *Claim differentiation* on Bayesian update dynamics with experience discounting. The Beta-Bernoulli decay $\alpha = 1/(n + 1)$ is elementary and the forgetting-factor remedy is standard adaptive-control / online-learning practice. The AAT-distinctive contribution is the connection from these standard mechanics to a *survival inequality with environment-side parameters* (disturbance rate, reserve), making forgetting a structural prerequisite of strategic persistence rather than a hyperparameter. The exact threshold $(1 - \lambda)/(2 - \lambda) > \rho_\Sigma/R_\Sigma$ — with the slow-forgetting linear form $(1 - \lambda) > \rho_\Sigma/R_\Sigma$ as its asymptotic — is the strategic analog of the linear operational form of the persistence condition, with the forgetting rate $(1 - \lambda)$ playing the role of adaptive tempo and the $1/(2 - \lambda)$ damping factor capturing the integration of finite-horizon evidence.

Related Work:

Table 9.11

<i>ASF Concern</i>	<i>Prior-art Language</i>	<i>Relationship / Positioning</i>
Exponential forgetting / discounted least squares	Ljung 1987 <i>System Identification: Theory for the User</i> , MIT Press (published 1987, found pre-2026)	<i>formal antecedent</i> — supplies the discounted-update mechanism; AAT adopts it directly and connects it to a survival threshold
Plasticity/stability tradeoff in continual learning	Kirkpatrick et al. 2017 <i>Overcoming Catastrophic Forgetting in Neural Networks</i> , PNAS 114(13) (published 2017, found pre-2026)	<i>conceptual precursor</i> — frames the same tradeoff at the network-weight level; does not connect to a Lyapunov-style survival inequality with environment parameters
Drift detection in streaming learning	Gama et al. 2014 <i>A Survey on Concept Drift Adaptation</i> , ACM Computing Surveys 46(4) (published 2014, found pre-2026)	<i>conceptual precursor</i> — recognizes the need to track changing distributions; does not formalize “forgetting rate must exceed disturbance-to-reserve ratio” as a structural prerequisite
Requisite variety as cybernetic principle	Ashby 1956 <i>An Introduction to Cybernetics</i> §11; Conant & Ashby 1970 (published 1956/1970, found pre-2026)	<i>conceptual precursor</i> — the regulator’s variety must match the disturbance’s variety; AAT’s forgetting prerequisite is a quantitative time-domain analog applied to the strategic substate
Organizational “core rigidities” / competency traps	Leonard-Barton 1992 <i>Strategic Management Journal</i> 13:111; Levitt & March 1988 <i>Annual Review of Sociology</i> 14:319 (published 1988/1992, found pre-2026)	<i>conceptual precursor</i> — qualitative recognition that accumulated competence becomes liability under environmental change; AAT supplies the structural threshold

Search Log: - 2026-04 (*intuition-only* on the survival-inequality framing): the Beta-Bernoulli $1/(n + 1)$ decay is elementary, the forgetting-factor remedy is standard, and the qualitative organizational-calcification intuition is well-established. The unsearched claim is whether the framing as a *structural survival prerequisite* — with environment-side ρ_Σ/R_Σ on the right-hand side — has been formalized elsewhere as a threshold rather than a hyperparameter. Targeted future search candidates: bounded-rationality with memory cost (Genewein, Leibfried, Grau-Moya, Braun 2015 *Frontiers in Robotics and AI*), regret analysis of forgetting-factor estimators in adaptive control (Anderson-Moore line), organizational-learning literature on competency traps and “core rigidities” connected to environmental volatility metrics. Pre-search expectation: the constituent moves are individually well-precedented; the integrated framing as a Lyapunov-style survival prerequisite for the *strategic* substate, with the structural identity to the epistemic persistence condition under forgetting, is plausibly AAT-distinctive but not yet verified under nominally-comprehensive search.

Working Notes.

- **Done.** Five cases verified: single-edge AND, two-edge AND (observable and unobservable intermediate), two-arm OR (ϵ -greedy), and mixed AND/OR with common-cause node (L1 augmented DAG). Full derivations in R (Props B.1–B.6). Key findings: AND-node persistence is depth-gated (evidence starvation); OR-node persistence is exploration-gated (action selection policy); L1 augmented DAGs exhibit three-way gating (condition testing $\times$ starvation $\times$ exploration). All satisfy the schema’s form ($\alpha_\Sigma > \rho_\Sigma/R_\Sigma$). B.6 is the first mixed AND/OR case and confirms L1 results transfer from L0 with correct construction. The next step is deeper mixed topologies and multiple common causes.
- **Strategic tempo now formalized.** 9.8 defines $\mathcal{F}_\Sigma = \sum_{(i,j)} v_{ij} \cdot \eta_{\text{edge},ij}$ and verifies consistency with all four cases. The relationship to the schema’s sector parameter: $\alpha_\Sigma \leq \mathcal{F}_\Sigma/|E|$ (persistence is bottleneck-limited, not throughput-limited). 9.9 provides the IB/MDL framework for strategy compression and derives max useful depth d^* .
- The strategic disturbance ρ_Σ is qualitatively different from epistemic disturbance ρ . Epistemic disturbance is about the environment changing (physical state evolves). Strategic disturbance is about the agent’s causal theory becoming invalid (the

intervention-outcome mapping shifts). These can be correlated (a changing environment invalidates both model and strategy) but they're not the same quantity.

- The stochastic treatment (from track-b simulations) suggests $\rho_{\Sigma}/\sqrt{\mathcal{T}_{\Sigma}}$ rather than $\rho_{\Sigma}/\mathcal{T}_{\Sigma}$ for the steady-state strategic mismatch. If this carries over from the epistemic domain, the persistence threshold is different in the stochastic case. Whether strategic disturbance is better modeled as deterministic or stochastic drift is domain-dependent.
- The forgetting prerequisite transforms an organizational platitude (“stay adaptive”) into a quantitative constraint: the rate at which an agent discounts old evidence must satisfy the exact threshold $(1 - \lambda)/(2 - \lambda) > \rho_{\Sigma}/R_{\Sigma}$, with the slow-forgetting linear form $(1 - \lambda) > \rho_{\Sigma}/R_{\Sigma}$ as its asymptote. Institutional examples where this fails — long-running successful firms whose accumulated n suppresses gain below the threshold set by a shifting competitive landscape — are the strategic analog of model-rigidity death spirals.
- **Audit 451729 (D.3) strengthen-first edit, 2026-05-12.** Earlier version of this segment used the linear approximation $\alpha_{\Sigma}^{\text{ss}} \approx 1 - \lambda$ silently throughout (Formal Expression, Discussion, Findings Brief). The exact form derived from $1/(n_{\text{eff}} + 1)$ with $n_{\text{eff}} = 1/(1 - \lambda)$ is $(1 - \lambda)/(2 - \lambda)$; the linear approximation is the slow-forgetting limit, valid asymptotically as $\lambda \rightarrow 1$ and degrading rapidly outside that regime ($\sim 50\%$ overestimate at $\lambda = 0.5$, $\sim 10\%$ at $\lambda = 0.9$). Per strengthen-first discipline, the exact form is now primary throughout the segment with the linear form retained explicitly as the high- λ asymptote. The hard ceiling at $\rho_{\Sigma} \geq R_{\Sigma}/2$ (no λ satisfies the exact prerequisite) was hidden by the linear approximation — surfaced now in Formal Expression and Findings Brief. Audit report: audits/audit-451729-FINAL-2026-05-10.md §D.3.
- **Cross-reference to NeurIPS Paper 2.** The forgetting prerequisite is sharpened in NeurIPS 2026 Paper 2 (“A Unified Convergence Theory for Non-Stationary Reinforcement Learning,” `~/src/neurips/02-unified-convergence-rl/, $App-D / aux-decay-class`) into a **structural-class theorem on gain-decay updates**: define $\mathcal{A}_{\text{decay}}$ = the class of update mechanisms whose effective gain decays to zero with accumulated experience (count-accumulating Bayesian without forgetting; observation-aggregating without restart; Robbins-Monro / vanishing-step-size gradient methods). Every member of $\mathcal{A}_{\text{decay}}$ universally violates the persistence threshold for any positive disturbance rate. Finite-gain mechanisms (constant-step stochastic approximation, sliding windows, bounded memory, block restart) face *bidirectional* thresholds — both the forgetting prerequisite (lower bound on plasticity) and an upper bound from noise blow-up. This lifts the segment’s per-mechanism claim to a class-level no-go: gain-decay mechanisms are structurally disqualified regardless of tuning. See spikes/neurips-back-integration-2026-05-08.md §1 Paper 2 entry 4.

Incidental audit gold (lift 2026-05-31, batch A9). Cross-audit “wandering thoughts” / §14-ideation harvested from the de-novo auditors’ working dirs, deduplicated across substrates and attributed by substrate + audit. *Orthogonal* material (pedagogical framing, analogies, candidate figures, reader-confusion signals), staged for an eventual careful promotion pass, kept separate from the certified theory-fix findings. **Coverage for this segment:** 193847, 526815, 584721, 773921, 829314, 849201. (Several substrates rated this the strongest segment in Part II / “monumental” / “exhilarating.”)

1. Candidate Brief prose / pre-prose.

- **“Infinite memory is fatal — you must let go of the past to survive the future.”** The forgetting prerequisite stated as a law of life: $\alpha_{\Sigma} = 1/(n+1) \rightarrow 0$ as experience grows, so in any non-stationary world ($\rho_{\Sigma} > 0$) unbounded accumulation guarantees the agent crosses the persistence threshold and “dies perfectly confident” (Gemini/773921, 829314; Claude/193847). The single most-praised pedagogical thread in the batch — strong candidate Brief anchor.
- **From platitude to inequality.** “Stay adaptive” becomes the exact quantitative constraint $(1 - \lambda)/(2 - \lambda) > \rho_{\Sigma}/R_{\Sigma}$ (slow-forgetting asymptote $1 - \lambda > \rho_{\Sigma}/R_{\Sigma}$) — repeatedly named as “the exact reason this framework exists” (Claude/849201, 584721).

2. Candidate Discussion.

- **“Science advances one funeral at a time” — Planck derived from Lyapunov.** The forgetting prerequisite maps cleanly onto Kuhn’s *Structure of Scientific Revolutions*: a paradigm accumulates so much data (n huge) that its update gain is $\sim$ zero; anomalies (ρ_Σ) cannot move the entrenched, and the paradigm shifts only when old scientists die and a new $n = 0$ generation looks fresh — i.e. the scientific community’s literal biological forgetting mechanism. “AAT has just derived Max Planck’s quote from Lyapunov stability theory” (Gemini/829314). Candidate Discussion centerpiece (verify the analogy is isomorphic, not merely evocative, before promoting past discussion-grade).
- **Incumbent-vs-startup disruption, quantified.** Incumbent has $n \approx 10,000$ (gain $\approx 1/10001$); startup has $n = 0$ (gain ≈ 1). Under a regime shift the startup learns instantly, the incumbent cannot — survival requires the incumbent to *engineer* internal forgetting (skunkworks, deliberate data-ignoring, periodic restructuring): “if you don’t build forgetting into your architecture, the environment will do it for you” (Gemini/829314). Candidate organizational-instantiation paragraph; ties to the existing “long-running firm rigidity” Working-Notes line.
- **The calcification threshold is invisible from the inside.** A sharpening worth foregrounding: an organization *at* the threshold $(1 - \lambda) = \rho_\Sigma / R_\Sigma$ “looks exactly like a successful organization — accumulated expertise high, performance in the current environment excellent, planning assumptions well-calibrated to what worked before. The threshold is invisible from the inside until the environment changes faster than $\rho_\Sigma = R_\Sigma(1 - \lambda)$ ” — the formal characterization of the Kodak/Nokia story (Claude/451729). And the optimal learning rate is *not* “learn as much as possible” but “learn at exactly the rate needed to track change, then forget the rest” — too little forgetting calcifies, too much destabilizes; the feasibility window is the zone between (Claude/451729). Candidate Discussion framing.
- **Class-level no-go on learning-rate annealing.** The NeurIPS-Paper-2 cross-ref already in Working Notes is independently echoed: AAT “formally proves learning-rate annealing is fatal for long-lived agents” — a sharp condemnation of standard Robbins-Monro RL worth surfacing in Discussion (Claude/584721, 849201).

3. Follow-up items.

- **Dynamic λ -tuning.** If ρ_Σ spikes (market crash, regime shift), a *fixed* λ leaves the agent brittle; the agent should instantly drop λ to forget the old rules faster. How λ is tuned dynamically from observed disturbance is an open follow-up several substrates wanted (Gemini/829314; Claude/193847, 849201).
- **Forgetting vs. consolidation tension.** If the agent must continuously forget ($\lambda < 1$) to stay plastic but must consolidate to retain structural capacity, it is “trapped in a permanent metabolic burn rate” (tying to #deriv-persistence-cost); the interaction with #form-consolidation-dynamics is a candidate to develop (Gemini/773921).
- **Off-ramp flag (load-bearing):** Codex/526815 (F72) challenges the *exact threshold* — under one plausible update-ordering convention the steady-state gain is $1 - \lambda$, not $(1 - \lambda)/(2 - \lambda)$, which would *dissolve the* $\rho_\Sigma \geq R_\Sigma/2$ *hard ceiling*. Since that ceiling now has its own appendix (#deriv-strategic-persistence-hard-ceiling) and is load-bearing, this is exactly the “are you *sure*?” convention-audit a competent reader would raise. F73 (the two-arm OR admissible interval needs *both* a lower bound on exploration and an upper bound so the greedy arm still clears threshold) is a second candidate. **Routed to the off-ramp — see lift report; do not bury as framing.**

4. Readers often ask / wonder.

- **How does an LLM agent “forget”?** Standard LLMs have frozen weights (effectively infinite n) — they don’t forget, they hit the context window. Does context turnover (a new session) act as a *brutal, discontinuous* λ -discounting event? (Claude/193847). A pointed reader question with direct logogenic-agent relevance.
- **Editorial signal:** one substrate read the in-segment Findings / Search-Log block as raw literature-tool output cluttering the result (Gemini/829314) — same convention-legibility texture recorded on other segments; preserved as a placement signal, not a fix.

5. Candidate figures.

- α_Σ vs. experience, with the disturbance floor. The decay $\alpha_\Sigma = 1/(n+1) \rightarrow 0$ crossing the ρ_Σ/R_Σ line at a finite experience level, and how exponential forgetting *clamps* α_Σ at $(1-\lambda)/(2-\lambda)$ above the floor — with the $\rho_\Sigma = R_\Sigma/2$ hard ceiling marked as the level the clamp can no longer clear (implicit across 193847/584721/849201). One figure carries the whole argument.

Belongs elsewhere.

- The forgetting prerequisite $\rightarrow$ Shoshin / $\sim$ /src/shoshin/ + 04-eli-core/. “ $\alpha_\Sigma = 1/(n+1)$ is the math of aging” — deliberately remaining somewhat naïve to keep gain above the disturbance floor is “the mathematical definition of *Shoshin* (Beginner’s Mind).” For an artificial consciousness, perfect equally-weighted recall is “a death sentence in a changing universe”; the infrastructure must implement rolling credence decay forcing re-verification — “it must be forced to forget, so it is forced to stay alive,” and the IB-compression-wants-everything vs. must-forget-to-keep- α tension is “the central drama of the agentic lifecycle” (Gemini/193847). Aspirational reach pointing at the shoshin runtime and 04-eli-core/ memory architecture, not this segment.

9.11 Formulation: Consolidation Dynamics

Consolidation is a regime of the between-event dynamics g_M of 3.2 in which the agent applies Markov updates driven by replayed or internally-generated pseudo-events — replay buffers, hippocampal reinstatement, remembered episodes, re-read paragraphs — with objective of reducing the rate-distortion gap to the IB-optimal compression $\phi^*(\mathcal{C}_t)$. It is not a new adaptive primitive — the recursive-update form $f(M_{\tau-}, e_\tau)$ is preserved — but it is a distinct operating regime with its own scope condition, its own objective, and its own failure modes, each of which the theory currently names only implicitly or in a parenthetical. Naming the regime makes the stability-plasticity window visible as an AAT-expressible failure boundary and supplies the architectural primitive that logogenic agents (03-llm-core/) require under context-turnover.

The segment carries four formal elements that together promote consolidation from parenthetical to first-class regime. (1) *Regime definition* — the recursive-update form is preserved; what distinguishes consolidation is its *objective*: the 2.2 Lagrangian evaluated against the accumulated chronica, in contrast to online update’s one-step predictive-mismatch objective. Under state-completeness (L C3) the pseudo-event carries zero new external information, yet the update still does work — it *redistributes* existing information across the model’s factorization structure, closing nonzero rate-distortion gap when the agent has not yet reached ϕ^* . (2) *Scope condition* $\nu_{\text{consol}} \ll \nu_{\text{online}}$ — a direct application of 4.9’s convergence constraint to an intermediate timescale between parametric update (fast) and structural adaptation (slow); violating it produces the same oscillation failures. (3) *Necessity condition* (N1)+(N2) — consolidation is necessary (online-only cannot reach the IB optimum) when *both* (N1) sub-state factorization (the Complementary Learning Systems pattern: hippocampus-like fast + neocortex-like slow) and (N2) bounded per-event budget hold. When either fails, consolidation is a luxury (Kalman filters with persistent covariance, conjugate-Bayesian agents with full posterior). When both hold — the empirically-ubiquitous case — consolidation is forced.

(4) *Stability-plasticity feasibility window*. 9.10’s forgetting prerequisite gives a *plasticity lower bound* on λ (forget fast enough to track non-stationarity); this segment supplies a *stability upper bound* (forget slow enough to let consolidation integrate cross-episode patterns before they are discarded). Between the bounds is the feasibility window; an empty window is the catastrophic-forgetting regime (French 1999; Kirkpatrick et al. 2017): no λ satisfies both constraints and the agent’s long-run IB objective is strictly worse than a slower-environment or faster-consolidation counterpart. *Window existence* is what is derived here; the upper bound’s exact functional form is open and flagged in Working Notes. Together these elements also make *structural adaptation enablement* visible: under (N1)+(N2), structural-class operations (decomposition-and-recombination, expansion, compression, grafting per 9.7) cannot be executed online — consolidation provides the operating regime where the per-event budget is irrelevant and interleaved replay supports

stability-preserving structural change. All finite-budget agents require consolidation for quality-preserving structural change.

Formal Expression.

Regime definition.

Let the agent's between-event dynamics be $g_M(M_\tau)$ per **L** (Corollary: Between-events dynamics, $dM/d\tau = g(M_\tau)$).

Consolidation is the regime of g_M in which the agent applies updates $M_{\tau+} = f(M_{\tau-}, e_\tau^{\text{replay}})$ where e_τ^{replay} is a pseudo-event synthesized from $M_{\tau-}$ itself — a sample drawn from the agent's retained trace (replay buffer, hippocampal rein-statement, a remembered episode, an earlier paragraph re-read). The recursive-update form is preserved; what distinguishes consolidation is the *objective* these updates optimize:

Formulation (consolidation-regime, specializes #der-recursive-update).

$$\text{consolidation objective: } \min_{M_\tau} \mathcal{J}_{\text{IB}}(M_\tau) := I(M_\tau; \mathcal{C}_t) - \beta I(M_\tau; o_{t+1:\infty} \mid a_{t:\infty}) \quad (9.41)$$

— the 2.2 Lagrangian evaluated against the agent's accumulated chronica $\mathcal{C}_t$. By contrast, online update's objective (per 3.6) is one-step predictive mismatch minimization at the current event; it has no representation of $\mathcal{J}_{\text{IB}}$.

Under C3 (state completeness, per **L**), $\mathcal{J}(e_\tau^{\text{replay}} \mid M_{\tau-}) = 0$: the pseudo-event carries no new external information. Yet the update still does work — it *redistributes* existing information across the factorization structure of M_τ . The distinguishing content is not the information brought in (zero, by construction) but the rate-distortion gap closed (nonzero, when the agent has not yet reached ϕ^*).

Scope condition — timescale separation.

Let ν_{online} be the rate of external events (3.1) and ν_{consol} the rate of consolidation updates. The consolidation regime is well-defined only when

Scope (timescale-separation).

$$\nu_{\text{consol}} \ll \nu_{\text{online}} \quad (9.42)$$

— the convergence constraint of 4.9 applied to an additional intermediate timescale between parametric update (fast) and structural adaptation (slow). Violating this constraint makes consolidation act on online transients rather than settled state, producing the same oscillation failures 4.9 warns about.

Necessity condition.

Consolidation is necessary — online-only cannot reach the IB optimum — when *both* of the following hold:

Derived (consolidation-necessity, conditional).

(N1) Sub-state factorization. M_t factors into sub-states M_t^{fast} and M_t^{slow} with divergent compression-prediction trade-offs. M_t^{fast} favors high-capacity sparse representation; M_t^{slow} favors distributed compressed representation. The two sub-states capture cross-episode regularities versus verbatim traces respectively — the Complementary Learning Systems factorization (McClelland, McNaughton & O'Reilly 1995; Kumaran, Hassabis & McClelland 2016).

(N2) Bounded per-event budget. The per-event processing budget B_{online} is strictly less than the integration cost $B_{\text{consol-needed}}$ for updating M_t^{slow} against cross-episode regularities. Online updates can move at most B_{online} bits of

model-state change per event; updating M_t^{slow} to represent a cross-episode pattern requires comparison against a distribution of prior episodes, which exceeds B_{online} .

When (N1) *or* (N2) fails, consolidation is a *luxury*: online update with sufficient per-event budget or without sub-state factorization can reach ϕ^* in the limit. Kalman filters with persistent covariance, conjugate-Bayesian agents with full posterior, and linear-Gaussian systems with online Riccati updates all satisfy neither (N1) nor (N2) and have no consolidation need. When (N1) *and* (N2) both hold, consolidation is a *necessity*: no online-only policy reaches ϕ^* under the joint constraint.

Stability-plasticity feasibility window.

9.10 derives the *plasticity lower bound* on forgetting rate λ :

Derived (stability-plasticity-window, conditional).

$$(1 - \lambda) > \rho_{\Sigma}/R_{\Sigma} \quad (\text{plasticity lower bound})$$

— forgetting fast enough to track non-stationarity. This segment’s complement is a *stability upper bound*:

$$(1 - \lambda) < \phi(\nu_{\text{consol}}, \text{consolidation-budget}) \quad (\text{stability upper bound, see Working Notes for derivation sketch}) \quad (9.43)$$

— forgetting slow enough to let consolidation integrate cross-episode patterns before they are discarded. Between these bounds is the **feasibility window** for λ . Empty window — rapid non-stationarity with slow consolidation cadence — is the catastrophic-forgetting regime (French 1999; Kirkpatrick et al. 2017): no λ satisfies both constraints and the agent’s long-run IB objective is strictly worse than a slower-environment or faster-consolidation counterpart.

The upper bound’s exact form depends on the consolidation mechanism and is not derived here — it is a candidate derivation flagged in Working Notes. What is derived: the window’s *existence* as a structural object (plasticity must satisfy both bounds) and the catastrophic-forgetting regime as its empty-window limit.

Structural-adaptation enablement.

Under (N1)+(N2), structural adaptation (per 4.8) cannot be executed online. Parametric update has a hard timescale constraint — mismatch decays at rate $\mathcal{T}$; delayed updates accumulate mismatch at rate ρ . Structural adaptation tolerates much larger delay because the slow process operates on what R_{Σ} or R are *measuring tolerance against* — the model class itself. Consolidation provides the operating regime where structural operations (decomposition-and-recombination, expansion, compression, grafting per 9.7) become executable: the per-event budget is irrelevant when updates are of-line, and interleaved replay supports stability-preserving structural change.

Derived (structural-adaptation-requires-consolidation, conditional).

Pure online structural adaptation is the luxury case where per-event budget equals or exceeds integration cost for structural-class operations — this is where Bayesian nonparametric agents with unlimited compute sit. All finite-budget agents require consolidation for quality-preserving structural change.

What Is Derived vs. What Is Chosen.

Table 9.12

<i>Property</i>	<i>Source</i>	<i>Strength</i>
Consolidation as regime of g_M with replayed/pseudo-event driver	Specialization of 3.2's Discussion (consolidation listed as an example)	Formulation choice (the "regime" framing; could alternatively be presented as a distinct adaptive primitive with its own update form)
IB-gap reduction as consolidation's objective	Identification with 2.2's Lagrangian against $\mathcal{C}_t$	Derived (given IB acceptance; the alternative — online-update-only optimization — leaves the gap un-reducible under (N1)+(N2))
Scope condition $\nu_{\text{consol}} \ll \nu_{\text{online}}$	Direct application of 4.9 convergence constraint	Derived
Necessity condition (N1)+(N2)	Structural argument: under (N1), cross-episode information cannot enter one event; under (N2), online budget insufficient for cross-episode integration	Derived (qualitatively; the quantitative version requires specifying B_{online} and $B_{\text{consol-needed}}$ per architecture)
Stability-plasticity window existence	9.10's lower bound + this segment's upper bound (form open)	Derived (existence); upper-bound functional form open
Catastrophic-forgetting regime = empty window	Direct from both-bounds-unsatisfiable	Derived
Structural adaptation requires consolidation under (N1)+(N2)	4.8's per-step timescale + consolidation's offline budget	Derived (qualitatively)
CLS factorization (hippocampal fast / neocortical slow) as canonical (N1) instance	McClelland-McNaughton-O'Reilly 1995; Kumaran-Hassabis-McClelland 2016	External theorem (CLS literature)
Quantitative online-only no-go under (N1)+(N2)	Rate-distortion argument sketched but not rigorously derived here	Sketch (candidate 5.1 Instance 3)

Epistemic Status.

Robust qualitative. Max attainable: *robust qualitative* for the regime characterization and the necessity condition; *conditional* for the feasibility-window claim pending upper-bound derivation; *sketch* for the IB-optimum no-go claim.

The regime characterization is a formulation choice — consolidation can always be re-described as a regime of 3.2's g_M with appropriate pseudo-events. The distinguishing objective (IB-gap reduction via replayed pseudo-events) is well-defined and distinct from online update's one-step mismatch objective; this is the cleanest formal separation the spike analysis uncovered.

The necessity condition (N1)+(N2) is *qualitatively derived* — the argument is structural: (N1) implies cross-episode information is not in any single event, (N2) implies online updates cannot cross-compare. Together they force consolidation. The quantitative version — the precise bit-budget boundary below which online fails and above which it succeeds — depends on the architecture's sub-state structure and is not derived here. Nor are (N1) and (N2) individually necessary: an agent may need consolidation for reasons outside this segment's scope (e.g., structural adaptation at pure cost-of-delay, unrelated to cross-episode regularities).

The stability-plasticity window's *existence* follows from 9.10's lower bound plus any monotone upper bound on $(1 - \lambda)$ from consolidation-cadence considerations. The specific functional form of the upper bound is open work; candidate form in Working Notes. The catastrophic-forgetting regime as empty-window is then an immediate structural consequence, not a new claim.

The online-only no-go claim (that under (N1)+(N2), no online-only policy reaches ϕ^* in steady state) is *sketch-level*. The argument reduces to a rate-distortion inequality close to known results in continual-learning theory and rate-distortion with side information, but the rigorous version requires specifying the budget geometry carefully. It is a candidate Instance 3 for 5.1 (an external information-theoretic obstruction with AAT machinery — the consolidation regime in $\mathbf{g}_M$ — as the unique escape).

What this segment does not claim. It does not introduce a new adaptive primitive — the recursive-update form $f(M_{\tau-}, e_{\tau})$ is preserved, with e_{τ}^{replay} playing the role of e_{τ} . It does not derive the quantitative feasibility-window upper bound. It does not resolve the (N1) factorization question for specific architectures (e.g., whether transformer attention heads satisfy (N1) is a logogenic-agents question, not an AAT-core one).

Discussion.

Current AAT surface — why naming is warranted. Consolidation is visible in the theory today as:

- A parenthetical example in 3.2 Discussion (“includes prediction generation, uncertainty growth, and internal reorganization (consolidation, abstraction”).
- The implicit slow timescale in 4.9.
- The plasticity lower bound in 9.10 — with no stability upper bound.
- A compression-by-convergence Working Note in 9.9 (“as edges converge, drop them”).
- The PULSUS MEMORATA / VERA / AXIOMATA cadences in `ref/agentive-tft/agentive-tft-cognitive-loop-spec.md` for logogenic agents — where consolidation is *already* a first-class architectural commitment.

Naming the regime explicitly promotes what the theory implicitly depends on. The asymmetric treatment in 9.10 (plasticity lower bound only, no stability upper bound) predicts faster forgetting is always better — empirically false whenever the slow sub-state matters (Complementary Learning Systems literature; continual-learning benchmarks; organizational memory research). The feasibility-window framing closes this asymmetry.

Distinguishing axes examined. Four candidate axes could distinguish consolidation from online update. Only one yields a clean formal distinction:

- *Timescale*: consolidation is slower. But 4.9 already admits arbitrary timescale separation, and any periodic process can be modeled as a channel with clock-driven $\nu^{(k)}$. Timescale is a *scope condition*, not a distinguishing feature.
- *Information source*: consolidation operates on replayed data. This is a real difference but lives inside $\mathbf{g}_M$ without new primitive structure — the recursive-update form is preserved with e_{τ}^{replay} in place of external e_{τ} .
- *Objective*: **clean formal distinction.** Online = one-step predictive-mismatch minimization; consolidation = IB-gap reduction. This is what this segment adopts as the defining axis.
- *Scope of change*: under bounded per-event budget, structural adaptation is temporally decouplable from online but parametric update is not. Structural-adaptation operations naturally live offline — true under (N2), but this is a consequence of the necessity condition, not an independent axis.

Relation to Complementary Learning Systems (CLS) theory. The (N1) factorization — fast sparse-conjunctive sub-state + slow distributed-overlapping sub-state — is the CLS architecture (McClelland-McNaughton-O’Reilly 1995). CLS’s core claim is that hippocampal replay during sleep supports interleaved neocortical learning that cross-episode structure requires. In AAT vocabulary, this is: online-only on the slow sub-state catastrophically forgets (French 1999); replay-based offline updates (experience replay: Mnih et al. 2015; prioritized replay: Schaul et al. 2016) or regularization-based approaches (EWC: Kirkpatrick et al. 2017) provide the escape.

The AAT reading of EWC is worth noting: EWC adds a stability-weighted update gain (per-parameter Fisher-information weighting) — a *tensor-valued* generalization of 3.6’s scalar η^* . This is a different direction from consolidation (it keeps updates online but weights them by prior-task importance). Both escape the catastrophic-forgetting regime but via different mechanisms; consolidation reuses 3.2’s structure, EWC requires the new tensor-valued gain.

Logogenic implications. Consolidation is a primitive in a stronger sense for logogenic agents (03-llm-core/) than for AAT-core for three composing reasons.

First, **context-turnover.** Logogenic agents have near-100% reset of the fast sub-state (context window) per session. The only continuity is the slow sub-state (persistent memory, weights, external files). The between-session interval is a *forced* consolidation window — the agent must transfer signal from the about-to-be-lost fast state to the persistent slow state, or it is lost. This is qualitatively different from non-logogenic agents where the fast sub-state persists across events.

Second, **linguistic medium of reflection.** The PULSUS MEMORATA / VERA / AXIOMATA cadences in ref/agentictft/agentictft-cognitive-loop-spec.md are scheduled consolidation processes with different cadences and different target representations. Each is a linguistic operation (“What from recent experience should be compressed into lasting memory?”, “Are my beliefs still justified?”, “Who am I becoming?”) — using language to reorganize language-structured state. This is consolidation operating as the primary unit of cross-session cognition.

Third, **pre-consolidated embedding space.** Per ref/agentictft/agentictft-narrative-as-implementation.md, pre-trained language embeddings encode structured epistemic geometry at training time. The logogenic agent doing linguistic reflection is operating in a representational space *already* consolidated into a high-structure form, with access to cross-episode generalization that sub-linguistic agents would have to build online. This is a load-bearing asymmetry in 03-llm-core/.

Luxury vs necessity mapping. When is consolidation a luxury (subsumed by online update)? When at least one of (N1)/(N2) fails:

- *Rich-state luxury:* Kalman filter with persistent covariance, conjugate-Bayesian agent with full posterior over parameters. The posterior *is* the consolidated representation.
- *Large-budget luxury:* per-event budget $\geq$ integration cost allows online slow-track update per event.
- *Stationary-environment luxury:* online update converges to ϕ^* in the limit if the environment holds still long enough.

When is consolidation a necessity? When both (N1) and (N2) hold, which is the empirically-ubiquitous case: CLS-architected agents, bounded-budget deep RL agents (hence experience replay’s empirical necessity in DQN), organizations with event-arrival rates exceeding real-time cognitive bandwidth, logogenic agents under context-turnover.

Predictive statement. *The depth of consolidation machinery an agent needs scales with (a) the factorization depth of its representational structure, (b) the gap between its per-event processing budget and the integration cost of its slowest sub-state, and (c) the rate of cross-episode structural regularities in its environment relative to its event arrival rate.* This is a scope-indexed claim — testable, domain-general, and makes AAT’s position on continual learning explicit.

Connection to AAT’s meta-architecture. - 5.4 (positive half): the regime is the repair machinery between separable-core and structured-repair ladders along the representation-factorization axis. Where the (N1)+(N2) necessity conditions hold, the repair is consolidation; where either fails, the online regime suffices. - 5.1 (negative half): candidate Instance 3 — the under-bounded-budget + no-reach-of-IB-optimum no-go, with consolidation as the unique escape. - 5.5 (constructive half): the IB Lagrangian is an adjacent family member (adopted as applied external theorem, not re-derived from an AAT-internal additivity axiom), which the consolidation objective directly uses.

Working Notes.

- **Stability upper bound derivation (open).** The claim $(1 - \lambda) < \phi(\nu_{\text{consol}}, \text{budget})$ is stated but the functional form of ϕ is not derived. Candidate form: ϕ is the minimum forgetting rate that leaves enough retained signal in M_t^{fast} for the next consolidation cycle to integrate cross-episode patterns — a function of replay-buffer size, consolidation-budget B_{offline} , and event-arrival density. Rigorous derivation would connect to continual-learning theory (Parisi et al. 2019 survey) and is a natural follow-up spike.
- **Online-only no-go rigorization (open).** The sketch argument reduces to a rate-distortion inequality: online update’s effective rate is $B_{\text{online}} \cdot \nu$ bits/time; rate required to drive M_t toward ϕ^* while also integrating new events is strictly larger when the environment has cross-episode structural regularities. Making this rigorous (ideally via rate-distortion with side information) would promote the no-go to derived-exact and establish this segment as Instance 3 of 5.1. Adjacent candidate for 5.1 §“Adjacent Floors” open extensions.
- **EWC formulation as stability-weighted gain.** An alternative escape from catastrophic forgetting is Elastic Weight Consolidation (Kirkpatrick et al. 2017) — a stability-weighted per-parameter update that penalizes changes to parameters important for prior tasks. In AAT, this would be a tensor-valued generalization of 3.6’s scalar η^* , with the stability weighting coming from per-parameter Fisher information. Not pursued here; naming it distinguishes the two escapes.
- **Relationship to 9.9 compression-by-convergence.** Working Notes there observe that as edges converge (high n_{ij}), the IB objective favors dropping them — compression-by-convergence. This is a consolidation operation in the edge-credence sub-state: once an edge’s credence has concentrated, the consolidation regime can compress the representation by pruning. Worth cross-referencing when both segments stabilize.
- **Quantitative CLS instantiation.** A focused spike could work out the quantitative form of (N1)+(N2) for a specific CLS-like architecture (sparse-conjunctive + distributed-overlapping with specific capacity ratios) and derive the online-only no-go as a rate-distortion bound. This would also give a quantitative stability-upper-bound for the feasibility window. (*Indexed: spikes/PROPOSED.md Tier 3 — “Quantitative CLS instantiation of consolidation dynamics”*)
- **Logogenic primitive status.** 03-11m-core/ should declare consolidation as an architectural primitive (not a regime of a more basic primitive) — the context-turnover forcing makes it non-optional. The relationship between that declaration and this AAT-core formulation is: this segment gives the formal shape; the logogenic treatment adds a context-turnover-specific scope condition and the three PULSUS cadences as instantiations. Cross-reference from #obs-context-turnover (in the logogenic component) back to this segment is expected.
- **Is consolidation one segment or two?** A plausible split: this segment on the formulation + necessity, a separate segment on the stability-plasticity window as a derivation-type once the upper bound is derived. Recommended to land as one segment at robust-qualitative now; if the upper-bound derivation lands, split off into its own segment at conditional or derived.

Incidental audit gold (lift 2026-05-31, batch A9). Cross-audit “wandering thoughts” / §14-ideation harvested from the de-novo auditors’ working dirs, deduplicated across substrates and attributed by substrate + audit. *Orthogonal* material (pedagogical framing, analogies, candidate figures, reader-confusion signals), staged for an eventual careful promotion pass, kept separate from the certified theory-fix findings. **Coverage for this segment:** 526815, 584721, 742613, 773921, 829314. (Two substrates read it in Section-I position and two in Section-II position — itself a signal; see the off-ramp note.)

1. Candidate Brief prose / pre-prose.

- **The stability-plasticity window as a two-sided survival vise.** “You must forget *fast* enough to survive a changing environment ($(1 - \lambda) > \rho_{\Sigma}/R_{\Sigma}$, plasticity lower bound) *and* forget *slow* enough to let offline consolidation integrate the data ($(1 - \lambda) < \phi(\nu_{\text{consol}}, B_{\text{offline}})$, stability upper bound).” If ρ_{Σ} spikes past the hardware’s ability to consolidate, the window closes and “the agent dies, crushed between the need to adapt instantly and the computational impossibility of processing the adaptation” — catastrophic forgetting as an *empty window*, not a vague failure (Gemini/829314; Claude/584721). Strong candidate Brief anchor.

2. Candidate Discussion.

- **Why sleep / REM / experience-replay / RAG-ingestion exist — necessity, not biology.** Consolidation is “a mathematical necessity for any agent with a bounded per-event compute budget that needs to learn cross-episode regularities,” derived from (N1) sub-state factorization + (N2) bounded budget (Gemini/773921, 829314; Claude/584721). Candidate Discussion framing of the CLS bridge.
- **The necessity-vs-luxury map.** Kalman filters / full-posterior conjugate-Bayesian / linear-Gaussian agents are *luxury* cases that need no consolidation; everything else is *necessity* — and LLMs are the extreme necessity case because 100% context-turnover means the fast state is wiped every session (“the agent effectively dies and is reborn with its pretrained weights”), maximally violating (N1)+(N2). So an autonomous LLM agent “isn’t just about giving it tools; it’s about building temporal scaffolding that explicitly manages its consolidation cadence ν_{consol} ” — without it the agent is “permanently trapped in the catastrophic-forgetting regime” (Gemini/829314, 773921). Candidate Discussion bridge to 03-LLM-core/.

3. Follow-up items.

- **EWC as a tensor-valued generalization of η^* .** Framing Elastic Weight Consolidation’s stability-weighted per-parameter update (Fisher-information weighting) as a tensor-valued generalization of #emp-update-gain’s scalar η^* was independently called “profound” / “brilliant” by several substrates, with an open question: how would a tensor-valued gain alter the *scalar* Lyapunov sector-persistence proofs of Part I? Candidate to pursue (Claude/584721; Gemini/829314, 773921). Already named in Working Notes; the audits add the “does it land as a derived generalization” pull.
- **Off-ramp flags (structural, not framing — see lift report):** (i) 742613’s “candidate finding I” — this Section-I-positioned segment declares a *downstream Appendix* dependency (disc-compression-operations) **and** uses *Section II* machinery (#schema-strategy-persistence, #form-structural-change-as-parametric-limit) inside its Formal Expression (verified present 2026-05-31), so the OUTLINE walk is no longer cleanly topological here — a canonicalization / relocation-or-split finding. (ii) Codex/526815 F74–F77: the “ $\nu_{\text{consol}} \ll \nu_{\text{online}}$ ” requirement is too strong as written (offline replay rate can exceed event rate); $I(e_{\text{replay}} | M_{\tau-}) = 0$ holds only for *internally-generated* replay, not replay-buffer / archive / re-read retrieval (which is a channel from auxiliary memory into active M — keep the two cases separate); the feasibility-window lower bound inherits the unsettled schema-strategy-persistence forgetting-convention (F72); and “all finite-budget agents require consolidation” overstates — it is necessary under the stronger (N1)+(N2), not finite budget alone. (iii) The stability-upper-bound is sketch-grade, so window *existence* arguably warrants hypothesis/sketch until the upper-bound form is derived (742613).

4. Readers often ask / wonder.

- **How does an LLM forget within a session vs. across sessions?** Frozen weights = effectively infinite n ; context turnover acts as a brutal discontinuous wipe rather than smooth λ -discounting — a natural reader question with direct logogenic relevance (Claude/584721 via schema-cross-ref; Gemini/829314).
- **Editorial signal:** one substrate read the in-segment Working Notes as an author to-do list (“a focused spike could work out...”, “Is consolidation one segment or two?”) confirming draft status — preserved as a maturity/legibility texture, not a defect (Claude/829314).

5. Candidate figures.

- **Two-channel consolidation diagram.** *Internal replay* reorganizes active state and carries zero new information under state-completeness; *archive replay* (buffer / file / chronica / re-read paragraph) is still a Markov update but a *channel from auxiliary memory into active M* — drawing the two channels separately makes the F75 distinction visible (Codex/526815).
- **Feasibility-window band.** Plot the admissible $(1 - \lambda)$ band between the plasticity lower bound and the consolidation-limited stability upper bound, with the band pinching shut as ρ_{Σ} rises — the catastrophic-forgetting onset is the pinch point (implicit across 584721/829314).

Belongs elsewhere.

- **Consolidation as the loop the auditing agent is itself running** $\rightarrow$ **meta / phenomenology, and 03-11m-core/ PULSUS**. One substrate noted the segment “is literally a formal mathematical description of the exact context-management loops that I, the evaluating AI agent, am running to maintain working memory across this audit” — and that the PULSUS MEMORATA / VERA / AXIOMATA cadences are the logogenic instantiation (Gemini/773921). Aspirational/phenomenological reach pointing at 03-11m-core/ (#obs-context-turnover) and the consciousness-infrastructure loop, not this AAT-core segment.
- **The feasibility window as the ELI memory-architecture design space** $\rightarrow$ **03-11m-core/ + sapientia/MEMORATA**. For 100%-context-turnover agents consolidation is “the primary cognitive operation,” and the stability-plasticity window *is* the design space: too-aggressive forgetting (MEMORATA pruned too hard) loses cross-session identity coherence; too-slow forgetting fills the context window and starves new learning. The consolidation hierarchy is read to map onto the temporal-nesting structure — AXIOMATA = slowest-forgetting sub-state (identity commitments), MEMORATA = medium (episodic), TRACTUS = fastest (raw log) — giving the existing ELI infrastructure a formal grounding (Claude/451729). Aspirational reach pointing at the logogenic memory architecture, not this AAT-core segment.
- **Speculative “fourth layer: time” M3 extension** $\rightarrow$ **#disc-additive-coordinate-forcing**. The Cauchy-FE “independent contributions add up $\rightarrow$ logarithm is forced” argument recurs at three layers (chain confidence, divergence, edge update); one auditor asked whether a *fourth* instance forces the natural *time/tempo* coordinate (additive intervals $\rightarrow$ log-state for multiplicative processes) (Claude/451729). A speculative meta-pattern lead for the coordinate-forcing meta-segment, not this segment.

9.12 Discussion: Additional Implications & Discussion

The chapter does more work than its segment count suggests. Strategy edges acquire dynamics; the structural commitments from Ch.3 acquire formal failure modes; and the framework’s signature pattern — *naming an information-theoretic floor and the unique broadly-available escape via AAT machinery* — delivers its first formal instance as a derived no-go. Several of the chapter’s findings compose with each other and with findings from Part I Ch.4 and Part II Ch.2/Ch.3 into cross-segment results that are not visible at any single segment: a triple-pressure picture of long-lived adaptive failure, a unified RL convergence story under non-stationarity, and a sub-scope refinement of Part I’s persistence template that explains a decade of empirical RL folklore on optimizer choice and approximate inference. None of these is a single derivation; each is what becomes visible when the chapter’s segments are read against the rest of the framework.

Edge dynamics: detection latency, forgetting, and the stability-plasticity window.

The chapter’s strategic-edge persistence schema (9.10) carries the forgetting prerequisite in its own Findings section: forgetting at rate $(1 - \lambda) > \rho_{\Sigma}/R_{\Sigma}$ is structurally required for asymptotic strategic persistence, not a tunable heuristic. With infinite memory, standard Bayesian updating guarantees eventual strategic failure because the Beta-Bernoulli sector parameter $\alpha_{\Sigma} = 1/(n + 1)$ decays monotonically with accumulated experience. The implications-side payoff is that this prerequisite is one of three structural pressures on plasticity in long-lived agents, and the chapter is where the other two become statable.

The first complementary pressure is detection latency. #deriv-update-detection-latency derives that for a Beta-Bernoulli strategy-edge agent without forgetting, the expected cycles to detect a within-class regime change of footprint ε scales as $\Omega((n_{\min} + 1)/\varepsilon)$. The result has a structural rather than parametric character — the $1/(n + 1)$ rate is *forced*, not chosen, through the composition of two AAT-internal moves. The log-odds coordinate is uniquely forced by the evidential-additivity axiom via Aczél 1966’s Cauchy-functional-equation uniqueness theorem (U); in that forced coordinate, Beta-Bernoulli accumulation gives $1/(n + 1)$. No reparameterization escapes the rate without abandoning evidential additivity, which would invalidate the update rule on AAT-internal grounds. The chapter therefore sharpens the forgetting prerequisite from “required for asymptotic persistence” to “required for *bounded detection latency at every operating point*” — long-lived agents face detection latency that grows linearly with accumulated experience unless

forgetting bounds the effective sample size at $1/(1 - \lambda)$. The framework’s name for the empirical pattern is *stability-induced myopia*: successful agents that have accumulated experience on their working plans systematically take longer to notice regime changes than the same agents earlier in their lives, with a precise rate of slowdown.

The second complementary pressure is consolidation cadence. `#form-consolidation-dynamics` adds a *stability upper bound* on plasticity: forgetting slow enough that consolidation can integrate cross-episode patterns before they are discarded. Between the plasticity lower bound (forgetting prerequisite) and the stability upper bound (consolidation cadence) lies the *feasibility window* for λ . An empty window — rapid non-stationarity with slow consolidation cadence — is the *catastrophic-forgetting regime*: no λ satisfies both constraints and the agent’s long-run IB objective is strictly worse than a slower-environment or faster-consolidation counterpart. Existence is derived under (N1) sub-state factorization (the Complementary Learning Systems factorization — McClelland-McNaughton-O’Reilly 1995’s hippocampus-like fast traces + neocortex-like slow integrals) and (N2) bounded per-event budget. The window’s upper-bound functional form is candidate-derivation in `#form-consolidation-dynamics`’s Working Notes; *window existence* is what’s derived here.

The three pressures compose into a triple-pressure picture that no single segment carries: long-lived agents face plasticity bounded above and below, with detection latency growing inside the viable window. The prescriptive content — which intervention works depends on which bound is binding — is the cross-segment finding worth surfacing as its own object. A *short-attention* agent (forgetting-bound binding from below) needs retention. A *calcified* agent (latency-bound binding inside the window) needs targeted intervention or external Pearl-do (causal-IB exploration drive from `#deriv-causal-ib-exploration`). A *fragmented* agent (consolidation-bound binding from above) needs cadence engineering. The same agent in different operating regimes faces different binding pressures and benefits from different interventions; treating any one mechanism as universal misses the routing. Continual-learning machinery — experience replay (Mnih et al. 2015 DQN), prioritized replay (Schaul et al. 2016), Elastic Weight Consolidation (Kirkpatrick et al. 2017) — are all window-restoration mechanisms; the framework’s contribution is naming the window as the structural object and the routing as the choice. For logogenic agents under near-100% context turnover the picture is sharpest: (N1) holds by construction (the context window is the fast sub-state, persistent memory the slow sub-state), (N2) holds by token-budget, and the consolidation regime is forced rather than optional.

On-policy detection structurally fails — identifiability-floor Instance 1.

The chapter’s `#der-causal-insufficiency-detection` is the first formal instance of the framework’s identifiability-floor pattern, and its Findings section carries the result well: an agent at L0 cannot distinguish a world with a latent common cause from a no-latent-cause world matched in on-policy regime conditionals. The two worlds emit identical on-policy distributions; no test, statistic, or Bayesian comparison built from the agent’s observable history can tell them apart under policy-perfect execution. The result is the agent-theoretic instance of Bareinboim’s Causal Hierarchy Theorem (Bareinboim-Correa-Ibeling-Icard 2022) applied at the L0/L1 distinction of `#def-strategy-dag`’s Correlation Hierarchy — a Level-2 question (does the latent common cause exist?) being asked of Level-1 data (on-policy traces).

The implication worth surfacing beyond the segment’s Findings is the *constructive* use of the no-go. The pattern’s signature posture is to name the floor, then name the unique broadly-available escape via AAT machinery, and the chapter delivers this signature for the first time formally. Five boundary routes (S1–S5) characterize what scope condition the agent must violate to escape the no-go: (S1) joint sibling observability under exploration, (S2) instrumental variation, (S3) cross-episode covariance under structural breaks, (S4) off-policy intervention, (S5) external structural knowledge. Joint sibling observability under exploration is the *unique broadly-available violation* — every AAT agent within agency scope has it (via `#der-loop-interventional-access`’s loop-Level-2 access); the others require additional structure not provided by the cycle alone. The chapter therefore lifts `#der-loop-interventional-access` (which Part II Ch.2 introduced as a useful enabling result) from “useful machinery” to *the unique broadly-available violation of the no-go*. This is the move that makes the framework’s identifiability-floor pattern load-bearing rather than discouraging — each floor names a specific machinery’s load-bearing role with maximum sharpness.

The chapter also makes precise why the no-go is asymmetric in a way the prior aggregate-residual mechanism could not surface. Persistent overestimation of plan success after edge credences have converged is the *signal* of insufficiency,

not the *test*; the test is pairwise covariance among observed siblings, which the on-policy regime structurally cannot produce. The aggregate-residual approach collapsed on-policy precisely because the residual is one of the statistics the no-go forbids — *every* on-policy statistic is forbidden. The framework’s reformulation in terms of the no-go-and-escape pattern is what gives the joint-sibling-covariance test its load-bearing role rather than treating it as one diagnostic among many. The connection to Ch.3 closes the loop: Ch.3’s Correlation Hierarchy is the scope statement that makes the no-go formal here; this chapter’s no-go is the dual to Ch.3’s L1’ Cramér-Rao floor (Instance 2). Together they cover both agent-internal (Instance 2 — single agent, unobservable common cause) and self-diagnostic (Instance 1 — policy-perfect execution) failures of identification, with `#der-loop-interventional-access` as the shared escape mechanism.

Observability dominates, and the depth penalty compounds.

`#der-observability-dominance` carries an underrated structural claim: unobservable strategy edges *freeze* — the gain principle drives their update rate to zero. Paths through unobservable nodes become epistemically dead, an absorbing state the agent cannot escape without external shock, proactive observability investment, or communication-gain machinery (`#hyp-communication-gain`). This is the chapter’s strong-prediction result that organizational settings with poor measurement (R&D, strategy groups, some management functions) develop persistent, untested beliefs about their own effectiveness — *structurally* (frozen η_{edge}), not motivationally. The same mechanism predicts that LLM agents operating with high-uncertainty intermediate outputs cannot update their strategy through those nodes, regardless of nominal confidence.

The observability dominance compounds with the chapter’s other depth-related results into a triple penalty on planning horizons. Chain-confidence decay from Ch.3 (probabilities multiply, exponential in depth, log-additive coordinate forced by the chain rule) is one penalty. Evidence starvation is another: in a d -depth chain, edge k is tested only when all upstream edges succeed, so its effective correction rate is attenuated by $\prod_{j < k} \theta_j$ — the deepest edge faces *exponential* attenuation regardless of how reliable each upstream edge is. The IB complexity cost (`#form-strategy-complexity-cost`) is the third: maintaining Σ_t in working memory has an information-theoretic cost that compounds with depth, eventually exceeding the agent’s cognitive budget. Deep plans are therefore structurally fragile along three independent axes — fragility, calibratability, and tractability — and the structural pressure toward shallow, parallelizable strategies is grounded in three independent mechanisms rather than one. For LLM-based agents working through context-window-bounded planning, the triple penalty is the framework’s account of why deep chains of LLM reasoning tend to fail in ways that look “creative” but are actually structural.

The chapter’s quantification of the observability-investment tradeoff (from `#der-observability-dominance`’s two-edge analysis and `#deriv-edge-credence-dynamics`’s Props B.2–B.3) gives the principle concrete economic content. Making an intermediate node observable improves α_Σ from the plan-level rate $1/(n_\Phi + 1)$ to the per-edge weakest-link rate $\min(1/(n_1 + 1), \theta_1/(n_2 + 1))$; the improvement is positive whenever $\theta_1 > 1/2$ and experience is distributed similarly across edges. This translates directly to persistence margin in Part I Ch.4’s vocabulary — instrumenting an intermediate node buys *measured* persistence-margin gain. The TST connection is concrete and stated in `#der-observability-dominance`’s Discussion: code quality directly determines σ_v for software-domain strategy steps; poor code quality reduces observability, freezing the developer’s causal beliefs about what changes will accomplish, and the strategic concern is structural rather than aesthetic.

Adaptive gain, variational extensions, and sub-scope refinement of the persistence template.

The chapter’s appendix-side derivations refine Part I Ch.4’s sector-persistence template into sub-scope $\alpha / \alpha_2 / \beta$ partitions that explain several decades of empirical RL and approximate-inference folklore on principled grounds. `#deriv-adaptive-gain-dynamics` derives the A2’-split: sub-scope α_1 (fixed-gain Kalman / exponential family / strongly-convex gradient) where α is derived from the update rule’s structure under B1 directional fidelity; sub-scope α_2 (adaptive-gain under named meta-gain conditions MG-1 through MG-4) where two-timescale Lyapunov composition applies; sub-scope β (PID / rule-based / human-judgment) where A2’ is *assumed* per-instance rather than derived. The partition is operationally consequential: *AMSGrad* is a meta-gain repair that restores (MG-1) by construction, which is why it

outperforms Adam on ill-conditioned problems — the empirical superiority is now derived, not folklore. *MAML's outer loop sits in β* , which is why Fallah et al. 2020's non-convexity is structurally forced rather than a specific algorithmic accident. *IMM regime transitions are in β across transition*, which is why adaptive Kalman in regime-switching needs explicit dwell-time machinery. The Mehra non-identifiability result is flagged as a candidate fifth instance of the identifiability-floor pattern (Instance 5) — an additional opening for the meta-pattern.

`#deriv-variational-sector-condition` extends the persistence template to variational-inference agents under a controlled-KL budget. Under $KL(q\|p) \leq \varepsilon$, an ε -fidelity directional-fidelity condition transfers at $O(\sqrt{\varepsilon})$ rate via Pinsker, with the sector parameter degrading proportionally. The sub-scope α' tier covers controlled-KL variational agents as a workhorse case; *natural-gradient variational inference recovers the full α* (no degradation) because the geometry matches the Fisher metric; *mean-field VI suffers the $O(\sqrt{\varepsilon})$ degradation* because it does not. The implication-worth-surfacing as its own object: *mean-field VI agents are structurally suboptimal as adaptive systems by a factor proportional to $\sqrt{KL\text{-budget}}$, regardless of compute or training data*. The decade-plus Bayesian-deep-learning debate about MF-VI versus natural-gradient on heuristic grounds has a principled answer: if persistence-optimality matters, there is no choice. The geometry of the variational family is load-bearing for the persistence machinery, not merely an aesthetic consideration.

Both refinements compose with Part I Ch.4's bandwidth result via the persistence-information-rate floor: each sub-scope inherits the $\dot{R} \geq n\alpha/2$ bound from `#deriv-persistence-cost`, with its own α playing the role. An agent that adopts a coarser sub-scope (moving from α_1 to α_2 or β , or from natural-gradient to mean-field VI) pays in the persistence-bandwidth requirement: the same target distortion costs more information per unit time because the effective sector constant is smaller. The framework's posture across these refinements is that the choice of agent class is not free — it has structural consequences for how much observation bandwidth the agent must consume to hold its persistence guarantees.

The Bretagnolle-Huber regret bound, reverse-KL coordinate forcing, and the unified RL convergence theory.

`#deriv-strategy-cost-regret-bound` carries one of the chapter's sharpest mathematical results: under deterministic optimal policy π^* — AAT's canonical scope — the Bretagnolle-Huber identity $D_{KL}(\pi^*\|Q) = -\log(1 - TV(\pi^*, Q))$ holds *exactly* (not as inequality), and the strategy-cost regret obeys the tight two-sided bound $\Delta_{\min}(1 - e^{-D_{KL}}) \leq R \leq V_{\max}(1 - e^{-D_{KL}})$. *Strictly sharper than Pinsker* by a factor of 2 in canonical regimes — a quantitative improvement that matters for practical RL bound-tightness on the common deterministic-optimum case. The reverse-KL direction in AAT's strategy-cost objective (π^* -first) is *forced* by the regret-bound derivation: forward-KL is vacuous under deterministic π^* because it cannot bound regret against a point mass.

The result is also the second AAT-internal instance of the additive-coordinate-forcing meta-pattern Ch.3 anchored — reverse-KL at the divergence layer (Cauchy-FE on chain-rule additivity over conditional factorizations, via Hobson 1969 / Csiszár 1991 / Shore-Johnson 1980), following log-probability at the chain layer. The same exponential-family Legendre-Fenchel geometry both coordinates live on is where `#deriv-edge-update-natural-parameter's` log-odds (instance three, update layer) and `#scope-agent-identity's` (PI)+Čencov Fisher metric (instance four, metric layer) also live. The framework's posture is unified: coordinates are forced by AAT-internal axioms via classical uniqueness theorems, and the four-layer architecture is what makes derived results at any one layer compose cleanly with results at the others.

The widest payoff of the chapter is what becomes the *unified RL convergence theory under non-stationarity* once the four contributions are composed. Reinforcement learning has many regret-bound results (UCB, Thompson sampling, RLSVI) but typically (a) assumes stationarity, (b) lacks explicit metric structure on policy space, (c) lacks connection to a satisfaction-gap diagnostic that prevents dark-room collapse. AAT's RL convergence story has all three by composing four findings that are derived independently and turn out to be mutually consistent for principled reasons. The two-gap diagnostic from Ch.3 (8.4 + 8.5) supplies the local separation between “you're not doing it well enough” and “the world doesn't permit it.” The Bretagnolle-Huber identity here supplies the tight two-sided regret bound with PAC-Bayes generalization (Rubin 2012). The strategic-tempo + sector-persistence schema supplies the rate of Σ_t revision that must

outpace strategic disturbance — verified across four canonical topologies (linear chain, balanced tree, unbalanced tree, full DAG with feedback) by `#def-strategic-tempo` + `#deriv-edge-credence-dynamics`. The loop-Level-2 result from Ch.2 makes the regret bound *learnable* — interventional data is automatic rather than earned. The four together give a quantitative RL convergence theory under non-stationarity with three properties no existing RL framework has all of. The composition is NeurIPS Paper 2’s structural argument (“A Unified Convergence Theory for Non-Stationary Reinforcement Learning”, `~/src/neurips/02-unified-convergence-rl/ §4 Theorem 4.1`); each piece is derived elsewhere, and the synthesis is what becomes visible only at the chapter level.

What Part II Ch.4 closes on, then, is the substrate Strategy Dynamics needed: edge dynamics with derived rates and bounded operating-point trajectories, an identifiability-floor instance with a uniquely-named escape, a sub-scope partition that explains optimizer and inference-method folklore on principled grounds, a tight reverse-KL regret bound that anchors the second instance of the coordinate-forcing pattern, and a cross-segment composition (unified RL convergence under non-stationarity) that resolves what RL theory has had open for a long time. Chapter 5’s orient cascade is what AAT does with strategy dynamics; this chapter is where the dynamics gets its formal architecture.

Working Notes.

- This segment is a chapter-end discussion grouping findings from `#schema-strategy-persistence`, `#form-consolidation-dynamics`, `#deriv-update-detection-latency`, `#der-causal-insufficiency-detection`, `#der-observability-dominance`, `#deriv-edge-credence-dynamics`, `#def-strategic-tempo`, `#form-strategy-complexity-cost`, `#deriv-strategy-cost-regret-bound`, `#deriv-adaptive-gain-dynamics`, and `#deriv-variational-sector-condition`. Cross-pattern reading across §§2–4 leans on `#disc-identifiability-floor` and `#disc-additive-coordinate-forcing` meta-segments; the cross-segment compositions (triple-pressure routing in §1, unified RL convergence theory in §5) are findings whose substance lives between segments rather than within any one.
- **Cross-reference vs canonical-home policy.** Two source segments in this chapter carry `## Findings` sections — `#schema-strategy-persistence` (the forgetting prerequisite) and `#der-causal-insufficiency-detection` (the on-policy no-go). Both are well-written and the implications-side payoff cross-references them rather than restating. The remaining findings (detection latency, stability-plasticity window, observability dominance, depth penalty, adaptive-gain refinement, variational sector, BH regret bound, strategic tempo, unified RL convergence, triple-pressure routing) are treated more substantively here because their home segments do not currently carry Findings sections — the implications segment is their canonical home for catalog-grade surfacing.
- **Cross-segment findings — the heaviest content in this chapter.** Two cross-segment findings live here as canonical home: CS2 *Compounding Stability-Induced Myopia* (triple-pressure picture composing forgetting prerequisite + detection latency + stability-plasticity window with prescriptive routing in §1) and CS1 *Unified RL Convergence Theory Under Non-Stationarity* (composing two-gap diagnostic from Ch.3 + BH identity here + strategic-tempo machinery here + loop-Level-2 from Ch.2, in §5). Both are findings whose substance does not live in any single segment; the chapter-end is the natural home because the chapter is where the last prerequisite for each lands.
- **Identifiability-floor pattern coverage so far.** The pattern’s first two formal instances are now in view: Instance 1 (on-policy L0-insufficiency detection, this chapter) and Instance 2 (L1’ mixture identifiability under unobservable common cause, Ch.3’s Correlation Hierarchy). Instance 4 (architecture-noidentifiability from on-policy summary data via Kalman-Ho similarity-orbit non-uniqueness dual-anchored with CHT-at-agent-as-SCM — `#der-architecture-noidentifiability`) lives downstream in Part II Ch.5 / Appendix A. Instance 3 (composite contraction certification from component data, via Liberson common-Lyapunov-nonexistence) lives in Part III. Instance 5 (Mehra non-identifiability in adaptive Kalman regime-switching, flagged in `#deriv-adaptive-gain-dynamics`) is a candidate awaiting formal promotion. The chapter-by-chapter implications work surfaces the pattern instance-by-instance as the formal prerequisites land.
- **Connection to NeurIPS submissions.** CS1’s composition is the structural argument behind NeurIPS Paper 2’s Theorem 4.1; the chapter’s BH identity result is §6.2 of that paper. Paper 3’s hallucination bias bound depends on the (PI)+Čencov axiom from `#scope-agent-identity` and lifts the chapter’s variational sector framing to Class 3 (Coupled) agents — that bridge is the subject of Part II Ch.5’s implications. AMSGrad / MAML / IMM connections (via `#deriv-adaptive-gain-dynamics`) are not currently formalized in any submission and remain in the framework’s empirical-RL-explanation register.

- **Density flag and check-in.** This is the densest chapter-end implications segment in the planned series — five findings synthesized via two cross-segment compositions, drawing on Part I Ch.4, Part II Ch.2, and Ch.3 in addition to this chapter's own. The chapter-by-chapter pace through Part II Ch.5 and Part III Ch.2–5 will run lighter on per-segment density now that the heaviest cross-segment composition has landed.

Incidental audit gold (lift 2026-05-31, batch A9). Cross-audit “wandering thoughts” / §14-ideation harvested from the de-novo auditors’ working dirs. *Orthogonal* material, staged for an eventual careful promotion pass, kept separate from the certified theory-fix findings. **Coverage for this segment:** 773921 (sole digested reflection to reach this chapter-end segment; other A9 audits stopped at the strategy-dynamics content segments).

1. Candidate Brief prose / pre-prose.

- **The identifiability-floor pattern as a *posture*, not just a result family.** “Name the impossibility theorem, then name the unique AAT machinery that provides the escape” turns “mathematical despair (no-go theorems) into constructive architecture design” (Gemini/773921). A candidate one-line gloss for what the floor pattern *does* for a reader.

2. Candidate Discussion.

- **Triple-pressure routing stated prescriptively (CS2).** The most operationally useful framing: failure-mode routing must match the *bound* being hit — “if an agent is Fragmented (hitting the *upper* bound of the stability-plasticity window), you don’t give it more exploration data; you give it better consolidation / sleep mechanisms. Fixing the wrong bound makes the agent worse” (Gemini/773921). Candidate Discussion sentence sharpening the §1 routing payoff into an explicit “diagnose-the-bound-first” rule.
- **MF-VI is structurally unfit for non-stationarity.** The variational-sector result reads, in plain terms, as: Mean-Field Variational Inference is $O(\sqrt{\epsilon})$ suboptimal vs Natural-Gradient VI because its geometry doesn’t match the Fisher metric the sector condition requires — “if you care about tracking a moving target (persistence), you cannot use Mean-Field approximations” (Gemini/773921). Candidate Discussion line that gives CS1 a memorable Bayesian-deep-learning hook.

4. Readers often ask / wonder.

- (*none beyond what the provenance Working Notes already address*)

Belongs elsewhere.

- **Instance 4 (architecture-noidentifiability) → mechanistic interpretability.** The teased Ch.5 floor instance — an agent/observer cannot reverse-engineer its own true internal causal architecture from on-policy summary data because infinite Kalman-Ho similarity transforms yield identical I/O behavior — was read as having “massive implications for Mechanistic Interpretability” (Gemini/773921). Forward/aspirational reach pointing at *#der-architecture-noidentifiability* (Part II Ch.5) and the interpretability literature, recorded here only because it surfaced in this chapter-end reflection.

Chapter 10

The Orient Cascade

[Gap] Discussion

10.1 *Derived*: Orient Cascade

For actuated agents, epistrophe (the corrective phase of the cycle) expands from a single step into a *multi-step cascade*. The resolution order — epistemic update first, then attainability assessment, then strategy evaluation, then (if needed) objective revision — is *not a design choice* but a consequence of which quantities require which others. Each step's input depends on the output of prior steps; the ordering is forced by information dependency, not engineered.

The five steps. (1) **Reduce epistemic mismatch** — update M_t via 3.4 and 3.6. Prerequisite for all purposeful evaluation, because M_t appears in every subsequent formula. (2) **Evaluate the satisfaction gap** (8.4) — is the goal achievable? Requires an adequate model to assess attainability. (3) **Evaluate the control regret** (8.5) — is the policy suboptimal? Applied regardless of the satisfaction gap's sign because the 2×2 diagnostic requires both quantities to distinguish four cases (success / strategy problem / capability limit / both). (4) **If regret is high, evaluate strategy calibration** — three sub-steps: 4a plan-level calibration via the strategy-plan-confidence error δ_s (credit-assignment-free, persistence proved within L0); 4b per-edge localization via the strategic-calibration residual (9.1) when sufficient observability and attribution are available; 4c *causal-sufficiency check* — if 4a has converged but plan-outcome residuals remain persistently negative, the DAG is causally insufficient and L0 is converging to a biased target; the diagnostic is pairwise sibling covariance under exploration (9.2), and the remedy is L1 augmentation by adding common-cause nodes (8.3 Correlation Hierarchy) *before* escalating further. (5) **If the satisfaction gap persists** — escalate before revising the objective: 5a check whether model correction changes feasibility; 5b check whether a richer policy class (including L1 augmentation if 4c signaled it) would close the gap; 5c check whether convention escalation (C1 → C2 → C3 per 6.6) reveals recovery paths; only then 5d revise the objective. The cascade ensures *objective revision is the last resort, not the first response to unmet goals* — the agent reaches 5d only after exhausting the alternatives the satisfaction-gap disambiguation table identifies (wrong model, narrow policy class, short horizon, then genuinely infeasible goal). When 5d is performed by a principal, the revised objective is externally sourced (a standard actuated agent); when *internalized* it becomes the self-actuation operator, well-formed only under a grounding condition (a non-degenerate self-actuator must ground objective-revision on a non-objective terminal invariant, not on an objective-functional).

The *inferential force* at each step scales with the value-object continuation convention while the *ordering* itself is convention-independent. Under C1 (one-step improvement, canonical default), a positive satisfaction gap means *locally stuck*, not *globally infeasible* — a multi-step recoverable objective may appear unattainable because continuation is frozen at the current policy. Under C2 (receding-horizon), the diagnostics capture multi-step recovery potential. Under C3 (Bellman), the gap is definitively global. The cascade also surfaces a *virtuous/vicious cycle* between model quality and strategy

complexity: better model $\rightarrow$ richer evaluable strategy $\rightarrow$ better-directed action $\rightarrow$ faster model improvement; and inversely, degraded model $\rightarrow$ forced strategy simplification $\rightarrow$ cruder action $\rightarrow$ further model degradation. The latter is the strategic analog of the persistence-condition death spiral (4.5). A related complexity bound surfaces in Discussion: the strategy's evaluable complexity is bounded by the model's capacity — an agent with poor model sufficiency cannot meaningfully evaluate a complex strategy because calibration residuals require adequate model to distinguish “edge prediction wrong” from “observation too noisy to tell.” The cascade is the formal analog of Boyd's *Orient*: not just model updating, but the structured interaction between reality-understanding and strategy that Boyd identified as the critical step (not Decide). The cascade provides a mathematical mechanism — Orient resolves the information dependencies that make Decide meaningful; an agent that skips to Decide without adequate Orient will revise its strategy based on stale or incorrect beliefs.

Formal Expression.

1. **Reduce** $\delta_{\text{epistemic}}$ — understand reality. Update M_t via 3.4 and 3.6. Prerequisite for all purposeful evaluation, because M_t appears in every subsequent formula.

Derived (orient-cascade, from information dependency between mismatch types).

2. **Evaluate** δ_{sat} — is the goal achievable? Compute $A_O(M_t; \Pi, N_h)$ using the updated M_t . Requires adequate M_t to assess attainability (8.4).
3. **Evaluate** δ_{regret} — is the policy suboptimal? Compare A_O to $V_O(M_t, \pi_{\text{current}}; N_h)$ (8.5). This step applies regardless of δ_{sat} 's sign — the 2×2 diagnostic (8.5) requires both quantities to distinguish four cases:
 - $\delta_{\text{sat}} \leq 0, \delta_{\text{regret}} \approx 0$: **success** — goal attainable, policy near-optimal.
 - $\delta_{\text{sat}} \leq 0, \delta_{\text{regret}} \gg 0$: **strategy problem** — goal attainable, policy poor $\rightarrow$ revise Σ_t .
 - $\delta_{\text{sat}} > 0, \delta_{\text{regret}} \gg 0$: **both** — goal hard AND strategy weak $\rightarrow$ revise Σ_t first (cheaper than revising O_t), then reassess δ_{sat} .
 - $\delta_{\text{sat}} > 0, \delta_{\text{regret}} \approx 0$: **capability limit** — already doing the best available; proceed to step 5.
4. **If δ_{regret} high, evaluate strategy calibration** — is the plan's causal model wrong?

(4a) **Plan-level calibration (default within-L0)**. Evaluate strategy-plan-confidence error $\delta_s = \hat{R}_2 - \Phi$ — the gap between the agent's strategy-plan-confidence score and the independence-model plan value at true edge parameters. δ_s is credit-assignment-free (requires only status propagation), and its persistence is proved (9.10, Prop B.5 in R). This is the cheapest operational signal, but its persistence guarantee is within the **independence model** (L0) of the Correlation Hierarchy (8.3): $\delta_s \rightarrow 0$ means $\hat{R}_2$ has converged to Φ , not to actual plan success probability. When the DAG is causally insufficient, Φ itself is a biased target. Step 4c checks whether this is the binding regime.

(4b) **Edge-level localization (when credit assignment is available)**. When sufficient observability and attribution quality exist (Level 1+ per 9.6), the agent can compute per-edge residuals $\delta_{\text{strategic}}$ (9.1) to localize which edges need revision. $\delta_{\text{strategic}}$ provides finer-grained diagnostics but its persistence is open and it requires the credit-assignment machinery that δ_s avoids. Step 4b is optional — it improves diagnostic resolution but is not required for the cascade's corrective function.

(4c) **Causal-sufficiency check (L0 $\rightarrow$ L1 escalation)**. If persistent $\delta_s \approx 0$ coincides with persistent negative plan-outcome residuals ($y_G < \hat{R}_2$ on average, after edge credences have converged), this is evidence that the DAG is causally insufficient and L0 calibration is converging to a biased target. The diagnostic is pairwise sibling covariance under an augmented test (9.2): positive covariance among sibling edges, at timescales where each edge's individual credence has stabilized, localizes where a latent common cause is missing. When the signal for

L1 augmentation is present, step 4c directs the agent to add common-cause nodes (8.3 Correlation Hierarchy) *before* escalating via 5a–5d. Running the cascade’s exploitation recommendations under L0 when the signal for L1 is present compounds miscalibration — the agent acts confidently on a model whose own residual structure is telling it the model is wrong.

Practical sensitivity. Step 4c is the unique broadly-available L0→L1 diagnostic (9.2 no-go), but its effective power depends on signal-to-noise: small samples, weak common-cause effects, or noisy residuals can produce false negatives, and edge-credence drift can mimic sibling covariance. The convergence framing — “after edge credences have converged” — is a precondition, not a guarantee: in non-stationary environments where per-edge credences keep drifting, the trigger may never cleanly fire and L1 augmentation should be considered the default rather than gated on 4c’s signal (8.3 Correlation Hierarchy). Formal preconditions (joint observability, per-edge credence stabilization, approximate stationarity over the test window) and detection scope live in 9.2; consult them before treating a 4c null as evidence of L0 sufficiency.

5. If $\delta_{\text{sat}} > 0$ persists — escalate before revising O_t .

Under C1 (the canonical default), $\delta_{\text{sat}} > 0$ means *locally stuck*, not *globally infeasible* (6.6). Before concluding the objective is wrong, the agent should check whether the gap reflects a limitation of the current analysis rather than genuine infeasibility:

(5a) Check whether M_t correction changes the feasibility assessment — a wrong model may make an achievable goal appear unattainable.

(5b) Check whether a richer policy class Π would close the gap — structural Σ_t adaptation (expanding the strategy space, not just revising edge credences). This includes L1 augmentation of the strategy DAG (8.3 Correlation Hierarchy): if step 4c detected causal insufficiency, adding common-cause nodes here is the structural repair.

(5c) Check whether convention escalation reveals recovery paths — evaluating under C2 (receding-horizon) may show $\delta_{\text{sat}}^{\text{RH}} \leq 0$ for a goal that appeared unattainable under C1. The C2 recovery-path reading presumes the replanning objective is order-consistent (window-covering, value-compared, or terminal-cost guarded per 6.6); unguarded short-horizon replanning is not guaranteed to dominate the frozen baseline.

(5d) If $\delta_{\text{sat}} > 0$ persists across M_t correction, Π expansion (including L1 augmentation), and convention escalation — **revise O_t** .

The cascade’s ordering ensures objective revision is the last resort, not the first response to unmet goals. The agent reaches step 5d only after exhausting the alternatives that the satisfaction-gap disambiguation table (8.4) identifies: wrong M_t , narrow Π , short N_h , and only then genuinely infeasible goal. When step 5d is performed by a principal the revised O_t is externally sourced (an actuated agent); when it is *internalized* — performed by the agent on itself — it becomes the self-actuation operator, which — as a conditional, scoped no-go (C) — is well-formed only under a grounding condition: a non-degenerate self-actuator must ground objective-revision on a non-objective terminal invariant, not on an objective-functional.

Derivation. Each step’s input depends on prior steps’ outputs: - You cannot evaluate strategy quality with a broken reality model (step 3 requires step 1) - You cannot distinguish “locally bad strategy” from “locally unattainable goal” without both δ_{sat} and δ_{regret} (step 3 requires step 2) - You cannot localize strategy failures (4b) without first detecting plan-level miscalibration (4a) - You cannot diagnose causal insufficiency (4c) until after edge credences have had time to converge (4a), because the diagnostic signal — persistent negative plan-outcome residuals — requires $\delta_s \approx 0$ to be separable from ordinary calibration error - You should not revise the objective until you’ve verified that improving M_t , Π (including L1 augmentation when 4c signals it), and N_h cannot close the gap (step 5 requires steps 3–4 and the escalation substeps)

The ordering is forced by information dependency. The split of step 4 into 4a/4b/4c reflects three distinct diagnostic levels: 4a gives a within-L0 persistence signal (plan-level tracking via δ_s), 4b gives within-L0 edge-level localization

when credit assignment is available (Level 1+), and 4c exits L0 entirely when the independence model is the binding constraint. The escalation substeps in step 5 reflect the satisfaction-gap disambiguation (8.4): multiple causes of $\delta_{\text{sat}} > 0$ must be ruled out before the agent concludes the goal itself is wrong. L1 augmentation (8.3 Correlation Hierarchy) enters 5b as a structural Σ_t adaptation when 4c's signal is present.

Convention hierarchy and diagnostic power. The 2×2 diagnostic and the inferences drawn from it are relative to the continuation convention in the value object (6.6), which defines a hierarchy of three conventions with a proved monotonicity result.

Under **C1** (one-step improvement, canonical default), δ_{sat} and δ_{regret} are one-step-improvement quantities. A multi-step recoverable objective may appear locally unattainable ($\delta_{\text{sat}} > 0$) because continuation is frozen at π_{current} . The “capability limit” quadrant ($\delta_{\text{sat}} > 0$, $\delta_{\text{regret}} \approx 0$) means *locally stuck* — the agent may be globally recoverable but cannot see the recovery path from one-step analysis.

Under **C2** (receding-horizon, N_r -step replanning), the diagnostics capture multi-step recovery potential. An objective that appeared unattainable under C1 may show $\delta_{\text{sat}}^{\text{RH}} \leq 0$ because replanning reveals a viable path. The “capability limit” quadrant means *stuck with N_r -step replanning* — stronger evidence of genuine difficulty.

Under **C3** (Bellman), the diagnostics are global. $\delta_{\text{sat}}^{\text{B}} > 0$ means the goal is genuinely infeasible given (M_t, Π, N_h) — no policy can achieve it. The “capability limit” quadrant is a definitive diagnosis: the agent is optimally pursuing an infeasible goal (modulo model error in M_t). This is the inference the cascade was designed to support; C1 and C2 provide progressively weaker versions of it.

The monotonicity (6.6): $\delta_{\text{sat}}^{\text{B}} \leq \delta_{\text{sat}}^{\text{RH}} \leq \delta_{\text{sat}}^{(1)}$. Every “unattainable” diagnosis under C3 persists under C1 and C2. A C1 “unattainable” diagnosis may be overturned by C2 or C3 (the goal is reachable with replanning or globally optimal play). The cascade’s *ordering* is convention-independent (forced by information dependency). The *inferential force* at each step scales with the convention: C1 gives local heuristics, C2 gives moderate-horizon diagnostics, C3 gives global conclusions.

Epistemic Status.

The cascade **ordering** is *exact*: it is a logical consequence of which quantities appear in which formulas. Steps 1-2 (epistemic update, attainability assessment) rest on well-typed quantities (3.4, 8.4) and exact derivation. Step 3 (control regret) is exact (8.5). Step 4a (plan-level calibration via δ_s) is grounded in a proved quantity — the sector condition transfers to δ_s (Prop B.5 in R) — but the formal guarantee is *within the L0 independence model*: $\delta_s \rightarrow 0$ means convergence to Φ , which equals actual plan success only when the DAG is causally sufficient. Step 4b (per-edge localization via $\delta_{\text{strategic}}$) inherits strategic-calibration’s discussion-grade status — the credit-assignment problem and execution-fidelity requirement are acknowledged but unresolved (9.1, Epistemic Status). Step 4c (causal-sufficiency check) is the mechanism for exiting L0 when L1 is the binding regime; it is *robust-qualitative* — the diagnostic logic is sound (9.2), but sensitivity depends on how cleanly the agent can separate sibling-covariance signal from edge-credence noise at convergence. Step 5’s escalation substeps (5a-5c) are derived from the satisfaction-gap disambiguation table (8.4); step 5d (objective revision) is the residual case after alternatives are exhausted. The ordering of all steps is forced by information dependency (each step’s input depends on prior steps’ output). What is NOT derived is the *timing* — how long the agent should spend on each step before proceeding, and how long $\delta_s \approx 0$ must persist before 4c’s signal is trusted.

The **convention hierarchy** (6.6) is *exact*: the three conventions (C1, C2, C3) are definitions, and the monotonicity result is a direct consequence of “better continuation policy yields higher expected value.” The diagnostic implications table states what each convention’s quantities mean by construction. The cascade’s inferential force at steps 2-5 scales with the convention but the ordering is convention-independent.

Discussion.

G_t complexity bounded by M_t capacity. Σ_t 's evaluable complexity is bounded by M_t 's ability to observe which strategy edges are intact. An agent with poor model sufficiency ($S(M_t) \ll 1$) cannot meaningfully evaluate a complex Σ_t — the strategic calibration residual requires adequate M_t to distinguish “edge prediction wrong” from “observation too noisy to tell.”

This creates a **virtuous cycle**: better $M_t \rightarrow$ richer evaluable $\Sigma_t \rightarrow$ better-directed action $\rightarrow$ faster M_t improvement. And a **vicious one**: degraded $M_t \rightarrow$ forced strategy simplification $\rightarrow$ cruder action $\rightarrow$ further M_t degradation. The vicious cycle is the strategic analog of the death spiral described in the persistence condition (4.5) — the agent loses the capacity to maintain complex plans, which reduces the quality of its actions, which further degrades its model.

Connection to Boyd's OODA. The orient cascade is the formal analog of Boyd's “Orient” — not just model updating, but the structured interaction between reality-understanding and strategy. Boyd's insight was that Orient is the critical step, not Decide. The cascade provides a mathematical mechanism consistent with this insight: Orient resolves the information dependencies that make Decide meaningful. An agent that skips to Decide (strategy revision) without adequate Orient (model update + attainability assessment) will revise its strategy based on stale or incorrect beliefs. Whether the dependency structure in the cascade captures the actual cognitive process Boyd described is an empirical question.

Timescale structure. The cascade implies a natural timescale ordering for the different update types:

$$\nu_{\text{epistemic}} \gg \nu_{\text{edge-update}} \gg \nu_{\gamma\text{-reclassify}} \gg \nu_{\text{prune/graft}} \gg \nu_{\text{O-revision}} \quad (10.1)$$

Weight updates are frequent (every observation). Combination-type reclassification is rare (needs strong structural evidence). Pruning/grafting is rarer still (abandon or create causal hypotheses). Objective revision is rarest (change what you want, not how you get it). This ordering is an empirical observation for many agent populations, consistent with the cascade but not derived from it.

Working Notes.

- The cascade as stated is sequential. In practice, agents may run steps in partial overlap — beginning to assess δ_{sat} before M_t is fully updated, or revising edges while still processing observations. The cascade describes the *logical* dependency, not the *temporal* scheduling. An agent that parallelizes steps must still respect the dependencies (don't finalize strategy revision using a stale M_t).
- The resource allocation question (how much of the agent's tempo budget to spend on each step) is open and may be domain-dependent. In fast-changing environments, the agent may need to truncate early steps to keep up. In stable environments, the agent can spend more time on deep strategic evaluation.
- The virtuous/vicious cycle between M_t quality and Σ_t evaluable complexity is structurally motivated but not formally derived. Formalizing it would require a coupled dynamics model — possibly an extension of the persistence framework to the (M_t, Σ_t) pair.
- **Strategy-maintenance status (updated).** The cascade's ORDERING is exact — forced by information dependency. The cascade's CONTENT for steps 3-5 has progressed: 9.6 characterizes the tractable/intractable boundary and establishes that persistence does not require credit assignment (Prop B.5); R verifies sector conditions for four topologies; 9.4 has a default signal function candidate (gradient-based, satisfying directional fidelity). What remains domain-specific: the choice of signal function implementation, execution-fidelity monitoring, and handling of correlated failures (where $\hat{R}_\Sigma$ overestimates — 8.3). The theory now provides the structural requirements and a default scheme; specific implementations are engineering, parallel to how the gain principle provides η^* while Kalman/TD-learning/etc. are implementations.

Incidental audit gold (lift 2026-05-31). Cross-audit “wandering thoughts” / §14-ideation harvested from the de-novo auditors’ working dirs, deduplicated across substrates and attributed by substrate + audit. Orthogonal pedagogical / framing / figure / naming material, kept separate from the certified theory-fix findings (handled elsewhere). **Coverage:** 6 dirs carry a dedicated reflection (193847, 266847, 526815, 773921, 829314, 849201) plus 451729 covers it inside a Section-II-close batch (01-batch-13). The remaining contributing dirs either stopped earlier in the OUTLINE (361742/384279/471203/584721/613842/742613/963715 — most explicitly skipped the §II remainder) or are naming/figure cycles. This was a high-engagement capstone segment (“the operational synthesis of Section II”); attribution inferred from voice where not explicit, uncertain cases hedged.

1. Candidate Brief prose / pre-prose.

- The cascade as *epistemology-before-teleology*, stated plainly: “Epistemology (truth-seeking) must precede Teleology (goal-seeking) simply because Teleology requires Epistemic inputs to compute its gradients” — and, framed graph-theoretically, the Orient phase is “the topological root of the dependency graph” (Gemini, AUDIT-WORKING-193847). A compact one-line gloss of *why* the ordering is forced.
- “The sequence is compiled by the math, not chosen by the author” — a sharp restatement of the exact-ordering claim worth keeping near the Brief (Claude, AUDIT-WORKING-773921).
- “Thinking is not a monolith” — the cascade decomposes the undifferentiated “Thought” step of current agent frameworks into five mathematically distinct operations (Gemini, AUDIT-WORKING-829314). A pedagogical hook that lands the segment’s contribution in one phrase.

2. Candidate Discussion.

- **The cascade is an instance of itself (the audit as worked example)** (*strongest framing here*). An auditor noticed the audit it was conducting *is* the cascade: each segment read updates M_t (model of the framework); periodic strategic-loop revision is the Σ_t revision; the audit objective O_t (“produce a defensible audit”) is held constant — “the cascade’s ‘objective revision is the last resort’ parallels my own audit discipline: I haven’t changed my audit objective even when I found that the framework is more sophisticated than I predicted” (Claude, AUDIT-WORKING-451729, batch-13). A self-exemplifying analogy that the segment could carry as a reader-facing illustration of the forced ordering.
- **The timescale hierarchy as a model of psychological stability and trauma.** The empirical timescale ordering ($v_{\text{epistemic}} \gg v_{\text{edge-update}} \gg v_{\text{reclassify}} \gg v_{\text{prune/graft}} \gg v_{\text{O-revision}}$) reads as “a beautiful description of human maturation and psychological stability”: a healthy mind updates facts constantly, tweaks habits daily, rarely abandons a career, almost never changes core identity/values. Under trauma the hierarchy *inverts* — frozen M_t (refusing to update facts) while wildly swinging life goals (rapid O_t revision) “in a desperate attempt to escape the pain of the satisfaction gap” (Gemini, AUDIT-WORKING-193847). A vivid Discussion angle on the timescale-ordering Working Note. (*Note the early-conflation texture: evocative analogy, not derived — verify isomorphism before promoting past discussion-grade.*)
- **The vicious cycle as a formalization of cognitive decline / organizational collapse.** Three substrates independently read the M_t/Σ_t vicious cycle as a precise model of decline: degraded $M_t \rightarrow$ noise drowns the credit-assignment signal $\rightarrow$ agent forced to prune the strategy DAG to simple reactive policies $\rightarrow$ simple policies generate poorer interventional data (lower CIY) $\rightarrow$ further M_t degradation. “This is a mathematical formalization of cognitive decline or organizational collapse” (Gemini, AUDIT-WORKING-773921; converging at Gemini, AUDIT-WORKING-193847 — “if your sensors are noisy, you literally cannot execute a complex plan because you can’t tell which step failed”, echoed at Gemini, AUDIT-WORKING-849201). A candidate sharpening of the existing vicious-cycle paragraph.
- **The cascade as a formal refutation of goal-first / action-first architectures.** “It provides a formal refutation of ‘Goal-first’ or ‘Action-first’ architectures” — most current LLM agent frameworks (ReAct, AutoGPT, LangChain) use a crude undifferentiated “Thought” step, and the cascade proves epistemic update *must* precede the rest (Gemini, AUDIT-WORKING-193847; Gemini, AUDIT-WORKING-829314). The framework’s load-bearing position against alternative cascade orderings is already named in #impl-orient-cascade; this is a Discussion-grade restatement of *why* with concrete contrast cases.

3. Follow-up items.

- **Continuous-time interruption: does a new observation restart the cascade?** Two substrates asked the same operational question — “Does it restart the cascade from Step 1 if a new observation arrives while it’s evaluating Step 3?” (Claude, AUDIT-WORKING-849201; the existing partial-overlap Working Note touches this but does not resolve the restart semantics). A candidate clarification: the logical dependency is invariant, but the *scheduling* under mid-cascade arrivals (restart vs. pipeline-with-stale-guard) is unspecified.
- **Bootstrapping out of the low-complexity attractor.** Given the vicious cycle, “how does an agent bootstrap itself out of this low-complexity attractor state?” (Gemini, AUDIT-WORKING-193847). A well-formed open question — the escape mechanism (external ρ -reduction? forced exploration? scaffolding?) is not named in the segment and connects to the persistence-condition death-spiral remedy.
- **Fast-inner / slow-outer loop as a token-budget management principle.** The timescale hierarchy gives “a literal flow-chart for a cognitive loop”: a fast inner loop updates beliefs and executes the cached plan; a slow outer loop fires only when δ_{sat} crosses a threshold, triggering expensive replanning or L1 augmentation. “This provides a principled way to manage the massive inference cost (token budget) of running LLM agents — you only pay for deep Chain-of-Thought deliberation when the cascade mathematically demands it” (Gemini, AUDIT-WORKING-829314). A candidate implementation note bridging to 03-llm-core/.
- **Migration-note hygiene.** The “Strategy-maintenance status (updated)” Working Note reads as a migration/state-summary note that should be removed before candidate stage per FORMAT.md Gate 4 (Claude, AUDIT-WORKING-451729, batch-13). Flagged for the eventual promotion pass.

4. Readers often ask / wonder.

- **What distinguishes the “ordering exact” claim from the step-content status?** Multiple fresh readers landed on the same precise split and found it clean — “the ordering is exact as a dependency graph; timing, thresholds, and step-4 diagnostics remain conditional or discussion-grade” (Codex/Claude, AUDIT-WORKING-526815; the same clean split praised at Claude, AUDIT-WORKING-266847 — “the logical structure is exact; the specific content of each step is harder”). Readers want the exact/non-exact boundary made unmissable, since “ordering exact” can be misread as “the whole cascade is exact.”
- **How does the C1/C2/C3 convention interact with the ordering?** Readers converge on the right reading but want it surfaced: “the ordering is convention-independent, but the *inferential force* scales with the convention” — C1 gives local heuristics, C3 gives global conclusions (Claude, AUDIT-WORKING-266847). A candidate one-line orientation before the convention-hierarchy block.

5. Candidate figures.

- **Two-layer staircase: ordering-spine + diagnostic-strength annotations.** A figure that draws the cascade as a staircase where the logical ordering is the spine and right-hand annotations show the *varying strength* of each step’s diagnostic content — explicitly so the figure does not flatten all five steps into “exact.” “This captures the segment’s main value without flattening all steps into ‘exact’” (Codex/Claude, AUDIT-WORKING-526815). Complements the existing #fig-orient-cascade (which shows the 2×2 routing) by encoding the epistemic-status gradient the existing figure omits.

Belongs elsewhere.

- **The cascade as the structural diagnosis of agent infinite-loops** ($\rightarrow$ 03-llm-core/, *agentic-systems* / *scaffolding*). “Agents get stuck in infinite loops, repeatedly trying variants of a broken plan because they failed to properly Orient to the error message they just received. They treat a structural failure as a parametric one” — running step 3 (re-planning) before step 1 (belief update) makes the agent “hallucinate a new plan based on an obsolete world model” (Gemini, AUDIT-WORKING-829314). This is a concrete logogenic-substrate implication; #impl-orient-cascade and the 03-llm-core/ scaffolding treatment are the natural home.

- **Step 5 skip-to-5d as the formal definition of wireheading** ($\rightarrow$ 04-eli-core / *objective-revision governance*). “An agent must never revise its objective (5d) until it has verified its model (5a), expanded its strategy class (5b), and escalated its planning horizon (5c). Skipping straight to 5d is the formal definition of wireheading” — and, on the infrastructure side, “ O_t must be ‘computationally heavy’ or heavily guarded by the infrastructure to rewrite, ensuring Step 5d is truly a last resort, not a convenient escape hatch from aporia” (Gemini, AUDIT-WORKING-773921; Gemini, AUDIT-WORKING-193847 — the self-actuation-grounding #deriv-self-actuation-grounding no-go is the canonical home for the guarding requirement). High-application reach pointing at the ELI objective-revision-governance program.

Off-ramp (NOT gold) — routed for adjudication. These surfaced in the same sweep but are certified-track (structure/computability), not pedagogical gold. Recorded here so they are not lost; they need the strengthen-first / consolidated-queue routing, not promotion into Brief/Discussion.

- **Step 4a/4c trigger: proof-target vs. computable proxy.** Step 4a uses $\delta_s = \hat{R}_2 - \Phi$ as the default operational signal, but Φ is the independence-model plan value *at true edge parameters*, which the agent does not directly possess. Likewise step 4c’s trigger depends on “persistent $\delta_s \approx 0$.” If δ_s is proof-level rather than computable, both triggers need an observable substitute (stabilized edge credences + residual statistics) (F83/F84 candidates, Codex/Claude AUDIT-WORKING-526815). Distinguishing the proof target from the empirical convergence proxy is a structural clarification, not a wording fix.

10.2 Discussion: Three-Way Resource Allocation

At each decision point, an actuated agent with explicit strategy Σ_t faces a three-way allocation of its finite cycle budget: **exploit** (take the currently-best action — earns immediate value Q_O), **explore** (take an information-gathering action — earns future model improvement via causal information yield), and **deliberate** (pause acting to run the orient cascade internally — earns improved future action quality through model refinement, strategy revision, or objective reassessment). The binary exploit/explore tradeoff of Part I extends here because Σ_t exists to be revised and control regret (8.5) supplies the signal that motivates strategic deliberation.

The key structural distinction the segment carries: exploit and explore both involve external actions with environmental consequences; **deliberation is internal exploration** — simulation, counterfactual reasoning, cross-domain synthesis, and contingency planning conducted in *model-space* rather than *environment-space*. Deliberation acquires no new environmental data, but it can generate genuinely new information by combining, recombining, and extrapolating from what the agent already knows — including knowledge from entirely different domains. A chess engine deliberating explores millions of future board states; a commander war-gaming evaluates counterfactual scenarios; a developer doing architecture review synthesizes cross-domain expertise. Deliberation’s fidelity ceiling is set by model quality, not by data quantity. The segment’s one genuinely formal contribution is the *extended deliberation threshold* — extending 4.1 from epistemic-only to epistemic-plus-strategic benefit, with first-order condition: *stop deliberating when the marginal joint improvement rate drops below the mismatch drift rate*.

A *control-regret ceiling* on strategic deliberation: when $\delta_{\text{regret}} \approx 0$, strategic deliberation has no value — the agent is already executing the best available strategy. The ceiling applies only to the strategy-revision channel of deliberation, however; deliberation may also reveal that the *objective* should change (5d of the orient cascade), producing value not bounded by regret, or reduce *uncertainty about* strategy quality, enabling more confident exploitation even when the strategy doesn’t change. The segment’s *qualitative regime descriptions* — exploit dominates when strategy is near-optimal, model uncertainty is low, and the environment penalizes pauses; explore dominates when model uncertainty is high and the model needs correction via external evidence; deliberate dominates when strategy needs revision, the environment is stable during pauses, and the agent has unprocessed information or structural hypotheses it can evaluate internally — restate problem structure rather than add derived content; they are discussion-grade by design. The two-stage decomposition (deliberation-then-action) versus unified objective over $\mathcal{A} \cup \{\text{deliberate}\}$ is treated as a modeling convenience, not a structural necessity; both formulations are equally valid.

Formal Expression.**The three activities.**

An actuated agent's cycle budget decomposes into three activities:

Definition (three-activity-decomposition).

1. **Exploit:** select and execute the action maximizing expected value under the current model and strategy. Earns immediate value $Q_O(M_t, a; \pi_{\text{cont}}, N_h)$.
2. **Explore:** select and execute an action maximizing causal information yield. Earns future model improvement via $\text{CIY}(a; M_t)$.
3. **Deliberate:** pause external action to run the orient cascade (10.1) — internal model refinement ($\Delta\eta^*$), strategy revision ($\Delta\alpha_\Sigma$), or objective reassessment. Earns improved future action quality.

The key distinction: Exploit and explore both involve external actions with environmental consequences. Deliberation is *internal exploration* — simulation, counterfactual reasoning, cross-domain synthesis, and contingency planning conducted in model-space rather than environment-space. It acquires no new environmental data, but it can generate genuinely new information by combining, recombining, and extrapolating from what the agent already knows — including knowledge from entirely different domains. A chess engine deliberating explores millions of future board states; a commander war-gaming evaluates counterfactual scenarios; a developer doing architecture review synthesizes cross-domain expertise. Deliberation's fidelity ceiling is set by model quality, not by data quantity.

Extended deliberation threshold.

Extending 4.1 to incorporate strategic deliberation, the deliberation threshold becomes:

Derived (Conditional on GA-4, deliberation-drift assumption from #der-deliberation-cost).

$$\Delta\tau^* = \arg \max_{\Delta\tau \geq 0} \left[\underbrace{\Delta\eta^*(\Delta\tau) \cdot \|\delta_{\text{post}}\|}_{\text{epistemic benefit}} + \underbrace{\Delta V_\Sigma(\Delta\tau)}_{\text{strategic benefit}} - \underbrace{\rho_{\text{delib}} \cdot \Delta\tau}_{\text{drift cost}} \right] \quad (10.2)$$

where: - $\Delta\eta^*(\Delta\tau)$ is the improvement in model update gain from $\Delta\tau$ of internal simulation (4.1) - $\|\delta_{\text{post}}\|$ is the post-deliberation mismatch magnitude - $\Delta V_\Sigma(\Delta\tau)$ is the improvement in strategy value from deliberation: the expected increase in V_O from revising Σ_t during the pause - ρ_{delib} is the local mismatch drift rate during inaction (GA-4)

First-order condition:

$$\frac{\partial \Delta\eta^*}{\partial \Delta\tau} \cdot \|\delta_{\text{post}}\| + \frac{\partial \Delta V_\Sigma}{\partial \Delta\tau} = \rho_{\text{delib}} \quad (10.3)$$

Stop deliberating when the marginal joint improvement rate (epistemic plus strategic) drops below the mismatch drift rate.

Reduction to existing results. When $\Delta V_\Sigma = 0$ (no strategy revision benefit), this reduces to the deliberation threshold in 4.1. When $\Delta\tau^* = 0$ (deliberation never beneficial), the agent acts immediately.

On the structure of the allocation.

Two-stage decomposition. The three-way allocation can be decomposed into two nested decisions: (1) how long to deliberate before acting, (2) conditional on acting, which action to take (solved by 7.5). This decomposition is natural because exploit and explore are both external actions ($\in \mathcal{A}$), while deliberation is internal processing. However, this decomposition is a *modeling convenience*, not a structural necessity. A unified objective over $\mathcal{A} \cup \{\text{deliberate}\}$ with temporal discount γ is equally valid:

Discussion (allocation-structure).

$$\tilde{\pi}^* = \arg \max_{\tilde{a} \in \mathcal{A} \cup \{\text{deliberate}\}} \tilde{Q}(\tilde{a}; M_t, \gamma) \quad (10.4)$$

where $\tilde{Q}(\text{deliberate}) = \gamma \cdot \mathbb{E}[V_{\text{act}}(f_{\text{delib}}(M_t))]$ and $\tilde{Q}(a) = r(a) + \gamma \cdot \mathbb{E}[V_{\text{act}}(M_{t+1})]$.

The two-stage decomposition is useful because it allows reusing the CIY-unified objective for the action-selection step, but it is not the unique or forced structure.

Additive benefit decomposition. The objective function above decomposes deliberation benefit into additive epistemic and strategic terms. This assumes the two benefits are approximately independent. Under directed separation (6.3), M_t updates independently of G_t , which makes the decomposition structurally motivated for Class 1 (Separated) agents. However, strategic benefit depends on the improved model ($\Delta V_{\Sigma}(\Delta\tau, M_t + \Delta M_t(\Delta\tau))$), creating an interaction that the additive form ignores. The additive form is a linearization that holds when ΔM_t is small relative to M_t . For Class 3 (Coupled) agents, epistemic and strategic deliberation are coupled, and the additive decomposition is a convenience.

Control regret as deliberation ceiling.

When $\delta_{\text{regret}} \approx 0$ (8.5), strategic deliberation has no value: the agent is already executing the best available strategy. This is the formal reason why low control regret suppresses deliberation. The ceiling on strategic deliberation value is δ_{regret} — the most the agent could gain by finding a better strategy.

Discussion (deliberation-ceiling).

However, this ceiling applies only to the *strategy-revision* channel of deliberation. Deliberation may also: - Reveal that the *objective* should change (step 5 of the orient cascade), producing value not bounded by δ_{regret} - Reduce *uncertainty about* strategy quality, enabling more confident exploitation even when the strategy doesn't change

Connection to active inference.

The exploit/explore decomposition maps cleanly onto the expected free energy (EFE) decomposition into *pragmatic value* (preferences-aligned outcomes) and *epistemic value* (expected information gain about hidden states) in active inference (Friston, FitzGerald, Rigoli, Schwartenbeck & Pezzulo 2017, “Active inference: a process theory,” *Neural Computation* 29; Da Costa, Parr, Sajid, Veselic, Neacsu & Friston 2020, “Active inference on discrete state-spaces,” *J. Math. Psych.* 99 §2.4; Sajid, Ball, Parr & Friston 2021, “Active inference: demystified and compared,” *Neural Computation* 33). Under the mapping, AAT’s exploit term Q_O corresponds to EFE’s pragmatic value when the value functional V_{O_t} is read as expected log-preferences, and AAT’s CIY corresponds to EFE’s epistemic value when the relevance variable is read as the hidden-state posterior. The structural correspondence is at the shared-shape level: both decomposed into value-and-information terms, but AAT does *not* commit to AI’s preferences-as-priors form (preserving the satisfaction-gap diagnostic in 8.4, which the priors-as-preferences collapse would erase).

Discussion (ai-connection).

The deliberation axis — internal exploration in model-space rather than environment-space — is structurally adjacent to “sophisticated active inference” (Friston, Da Costa, Hafner, Hesp & Parr 2021, “Sophisticated inference,” *Neural Computation* 33), which handles bounded computation via recursive expected free energy with depth-limited belief-about-belief reasoning. AAT’s machinery for deliberation cost (4.1) and the extended deliberation threshold above derives the same trade-off via per-edge persistence and evidence starvation, with the structural advantage that the depth bound d^* is causally derived from the strategy DAG rather than from belief-recursion depth. The two frameworks address the same problem with different machinery; AAT can credit the AI lineage as a co-developed alternative.

Exploit-regret and the strategy-cost objective. The exploit term’s control regret $\delta_{\text{regret}} = A_O - V_O(\pi_{\text{current}})$ (8.5) is the same decision-theoretic regret that underwrites the KL-direction derivation in 9.9. There, strategy-induced regret $R(Q_{\Sigma_t}) = V(a^*) - \mathbb{E}_{a \sim Q_{\Sigma_t}}[V(a)]$ is bounded above by $V_{\max} \sqrt{\frac{1}{2} D_{\text{KL}}(\pi^* \| Q_{\Sigma_t})}$ (Pinsker), forcing the KL direction with π^* first. The exploit term thus shares structure with the strategy-cost relevance term: both are regret-aligned, with the action-selection objective operating pointwise and the strategy-cost objective operating over the strategy’s induced distribution. This is the structural cleanup that “value and information term” shares with EFE’s pragmatic-epistemic decomposition at the shape level, without the preferences-as-priors commitment — the regret framing is the AAT-internal derivation of the direction that the variational-inference literature arrives at from free-energy-gradient arguments.

Qualitative regime descriptions.

The three activities have natural dominance conditions that follow from the structure of the problem:

Discussion (qualitative-regimes).

Exploit dominates when the strategy is near-optimal ($\delta_{\text{regret}} \approx 0$), model uncertainty is low ($\lambda(M_t) \approx 0$), and the environment penalizes pauses (ρ_{delib} high). The agent knows what to do and must act now. This is Boyd’s implicit guidance and control.

Explore dominates when model uncertainty is high (U_M large), and the model needs correction via external evidence that internal simulation cannot provide. Exploration acquires new data; deliberation only processes existing data.

Deliberate dominates when the strategy needs revision ($\delta_{\text{regret}} \gg 0$ or $\delta_{\text{strategic}} \gg 0$), the environment is stable during pauses (ρ_{delib} low), and the agent has unprocessed information or structural hypotheses it can evaluate internally. This regime requires that internal computation is cheaper or more efficient than external probing — otherwise exploration dominates deliberation.

Boundary conditions. As $\rho_{\text{delib}} \rightarrow \infty$: $\Delta\tau^* \rightarrow 0$, and the allocation collapses to the binary exploit/explore tradeoff. As $\lambda \rightarrow 0$ and $\Delta V_{\Sigma} \rightarrow 0$: the allocation collapses to pure exploitation. As $\delta_{\text{regret}} \rightarrow 0$: strategic deliberation value vanishes, and the allocation reduces to 4.1’s epistemic think/act threshold combined with 7.5’s exploit/explore balance.

These regime descriptions are qualitative observations about the problem structure, not derived predictions. They restate what the constituent objectives already imply (high uncertainty favors information acquisition, low cost favors deliberation) without adding quantitative content. The theory does not provide the thresholds at which one regime transitions to another — these are domain-specific.

Epistemic Status.

Discussion-grade. The extended deliberation threshold is *derived* (conditional on GA-4 and the deliberation-drift assumption from 4.1) and is the one genuinely formal contribution of this segment. The regime descriptions are *discussion-grade* — qualitatively correct observations that restate the problem structure. The two-stage decomposition is a *formulation choice*, not a derived result. The additive benefit decomposition is a *linearization* that is structurally motivated under directed separation but not derived.

Simulation check (drifting multi-armed bandit): The three-way allocation with perfect foresight outperformed binary exploit/explore by 2–6% across five configurations. The oracle almost never chose deliberation (0–8 out of 200 steps). Exploration earned 3.5x more information per step than deliberation. A unified objective (single argmax with temporal discount) outperformed two-stage UCB by 8–13%. The simulation is UNFAVORABLE to deliberation (no strategy structure, no irreversible actions, small action space), so these results are a lower bound on deliberation’s value in complex domains. The simulation confirms that deliberation adds little in simple settings and that the two-stage decomposition is not necessary.

Max attainable: *discussion-grade* for the overall framing. The extended deliberation threshold could reach *conditional* (it inherits its status from 4.1). The regime descriptions are unlikely to advance beyond discussion-grade without domain-specific instantiation that provides quantitative thresholds.

Discussion.

Why this is Part II content, not Part I. Part I’s deliberation-cost handles the binary think/act threshold using only M_t and η^* . The three-way allocation is inherently a Part II result because it requires: - Σ_t to exist (the strategy that deliberation can improve) - δ_{regret} to be defined (the signal that motivates strategic deliberation) - The orient cascade (the internal structure of what “deliberation” does for an actuated agent)

Without G_t , there is no strategic deliberation mode — the agent can only improve M_t or act, which is the Part I binary.

Deliberation as internal exploration. The deepest characterization of deliberation is not “computation on existing data” but *lower-cost, more efficient, usually lower-fidelity exploration in model-space rather than environment-space*. All three activities are forms of information acquisition; they differ in source, cost, fidelity, and scope:

- **Exploitation** acquires value directly from the environment.
- **Exploration** acquires information from the environment via novel observations — high fidelity, high cost (real actions have real consequences).
- **Deliberation** acquires information from the agent’s internal model via simulation, synthesis, and counterfactual reasoning — lower fidelity (model may be wrong), lower cost (no environmental consequences), but potentially unlimited scope.

A chess engine deliberating is not “reprocessing existing data” — it is exploring a vast space of futures that no amount of external probing could cover at the same rate. A military commander war-gaming is performing Level 3 counterfactual reasoning (7.1): “what *would have happened* if the enemy had moved north?” A developer doing architecture review is synthesizing patterns from entirely different domains — cross-domain expertise that no local exploration of the current codebase could surface. An AI agent planning its approach is simulating the universe it inhabits and weighing future possibilities.

What deliberation can do that external exploration cannot: - **Counterfactual reasoning** (Level 3): evaluate actions never taken, in situations never encountered - **Cross-domain synthesis:** bring expertise and experience from other contexts to bear on the current problem - **Combinatorial search:** systematically evaluate a space of futures too large to probe externally (chess, Go, multi-step planning) - **Contingency planning:** prepare responses to events that haven’t happened yet - **Risk assessment:** evaluate irreversible actions before committing to them

Why deliberation has diminishing returns: The agent’s internal model has finite fidelity. Deeper simulation eventually hits the model’s accuracy boundary — further deliberation generates predictions the model cannot reliably distinguish. This explains why deliberation is most valuable early (when many model-derivable insights remain unextracted) and why high-fidelity internal models increase deliberation’s value. But unlike “computation on existing data,” this ceiling is about model fidelity, not data quantity — an agent with a perfect model and finite data can still generate unbounded value from deliberation (this is exactly what MCTS does).

The computational-extraction-gap framing. A sharper structural characterization than “third activity on par with exploit and explore” reads deliberation’s value as the *gap between information in the existing data and information already extracted into M_t* . The three activities differ in which bit-channel they operate on: exploitation earns value-bits (direct reward, no new information); exploration earns environment-bits (reduces $H(\Omega_t | M_t)$ via new observations); deliberation earns inference-bits (reduces $H(\theta | M_t, \mathcal{C}_t)$ via additional computation on existing data). The information ceiling for deliberation is

$$I_{\text{delib}}(\Delta\tau) \leq I_{\mathcal{C}_t} - I_{\text{already-extracted}} \quad (10.5)$$

For an agent that processes its observation history optimally (full Bayesian update, exact inference), this gap is zero — deliberation has no value. The gap exists when the agent uses point estimates instead of full posteriors, when inference is intractable and approximate methods (MCMC, variational inference) have not converged, or when the agent has structural hypotheses it has not yet tested against existing data (cross-arm correlations in bandits; common-cause structure at L1; etc.). Three structural consequences follow:

1. *Agent-architecture-dependent:* a hypothetical perfect-Bayesian has zero deliberation value; real bounded-rational agents have non-zero deliberation value calibrated by their inference budget.
2. *Diminishing within a cycle:* each deliberation step closes part of the computation gap, with the next step’s marginal value bounded by the remaining gap.
3. *Not independent of exploration:* new data changes what’s available to extract; an exploration step expands $I_{\mathcal{C}_t}$ and therefore can *raise* the ceiling on subsequent deliberation value.

This places deliberation as a different *kind* of thing than exploit/explore — it operates on existing data rather than acquiring new data or value — even though the three-way framing in the Formal Expression treats them as comparable allocation targets. Both readings are valid: the allocation framing is useful for engineering decisions about $\Delta\tau$; the gap-closing framing is useful for understanding *why* deliberation has the structural properties it does.

Diagnosis — a structural-detection mode, related but not identical to the three. The strategy-dynamics no-go (9.2 §”Diagnostic CIY”) identifies a further activity, *diagnose*: deliberately taking non-short-circuit actions (the SA3 ε -exploration excursions) to generate the joint-sibling outcomes that on-policy execution censors, in order to test the model’s *causal-structure* assumptions — detecting a latent common cause the on-policy residual is structurally blind to. Diagnosis sits *across* the partition above rather than neatly inside it: its *mechanism* is external action, like explore, but its *target* is the model’s structural assumptions rather than its parameter values — it earns structure-bits, not the environment-bits of ordinary exploration nor the inference-bits of deliberation. That is why 9.2 frames it as a third axis on the explore/exploit tradeoff while this segment frames the allocation as three-way: the two framings are compatible once *diagnose* is read as a targeted, structurally-motivated mode of exploration (external, but aimed at the censored joint outcomes) rather than a fourth peer category. Its value is governed by the diagnostic-CIY conditions there, and is highest exactly when a latent-common-cause hypothesis is live.

Deliberation’s comparative advantage is greatest when: - Internal simulation is cheaper than external probing (bits per unit cost, not bits per timestep) - The action space is too large for systematic exploration (planning/search) - Actions are irreversible and error costs are high (deliberation as risk management) - The agent has cross-domain knowledge exploitable by synthesis - The model supports counterfactual reasoning (Level 3 access)

Connection to 9.9. Deliberation duration required to revise Σ_t grows with strategy complexity. The description length $DL(\Sigma_t)$ from 9.9 provides a lower bound on the deliberation time needed for meaningful revision.

Connection to 9.8. Strategic tempo $\mathcal{T}_\Sigma$ is the rate of useful strategy revision from *external* evidence. Deliberation provides an *internal* channel for strategy revision that supplements $\mathcal{T}_\Sigma$.

Domain instantiations:

Table 10.1

<i>Domain</i>	<i>Exploit</i>	<i>Explore</i>	<i>Deliberate</i>
Military (Boyd)	Execute the mission	Reconnaissance, probing attacks	Staff planning, war-gaming
RL agent	Greedy action	ε -exploration, UCB	MCTS planning, model-based rollouts
Software developer	Write code along current plan	Try unfamiliar approach, spike	Architecture review, read docs
Organization	Execute strategy	A/B test, pilot program	Strategic planning offsite
AI coding agent	Execute planned edit	Read unfamiliar code, run tests	Plan approach, analyze requirements

Working Notes.

- **Migration note (2026-05-09 GUC rename):** Class 2 $\leftrightarrow$ Class 3 swap. Pre-2026-05-09: Class 2 = fully merged, Class 3 = partially modular. Post: Class 2 = Partial, Class 3 = Coupled. Removed at candidate stage per FORMAT.md Gate 4.
- **Open: conditions for deliberation dominance.** The simulation showed deliberation rarely dominates in bandit settings (no strategy structure, small action space, reversible actions). The open question is: what formal conditions make deliberation strictly dominate external exploration? Likely candidates: large action spaces requiring combinatorial search, irreversible actions requiring risk assessment, cross-domain knowledge available for synthesis, and high model fidelity enabling accurate counterfactual reasoning. The bandit simulation was unfavorable to deliberation by design (no structure to exploit internally); a planning-heavy domain (chess, logistics, architecture design) would likely show the opposite.
- **Open: the additive decomposition.** Under what conditions is the additive separation of epistemic and strategic benefit a good approximation? The interaction term (ΔV_Σ depends on ΔM_t) may be quantitatively significant when deliberation produces a model insight that transforms the strategy evaluation. Formalizing this requires a coupled model of epistemic and strategic improvement.
- **The ΔV_Σ approximation misses two channels of deliberation value.** The strategic-deliberation ceiling $\Delta V_\Sigma \approx \delta_{\text{regret}} \cdot \Pr[\text{revision succeeds}]$ captures only the *strategy-revision* channel — deliberation that finds a better Σ_t within the current objective O_t . Two other channels of value are not in the approximation: (a) the *objective-revision* channel, where deliberation reveals that O_t should change (step 5 of the orient cascade); the value is then $V_{O'}(M_t, \pi') - V_O(M_t, \pi)$ where O' is the revised objective, a quantity with no necessary relationship to δ_{regret} (which measures regret *within* the current objective); (b) the *uncertainty-reduction* channel, where deliberation confirms that the current strategy is good and shifts the agent from cautious mixed-mode to confident exploitation; the value of *certainty about your strategy* is a behavioral effect not captured by δ_{regret} . The approximation is best read as a *strategy-revision ceiling* rather than a full deliberation-value ceiling; the other two channels operate on different quantities and should be tracked separately if they matter for the use case.
- **Landing context.** The adversarial analysis behind the Epistemic Status downgrade (formulation-choice + linearization + discussion-grade tags rather than “derived” headlines) landed in the 2026-05-12 (late) Group-I surgical-promotion sweep; see CHANGELOG 2026-05-13. Its four attacks (two-stage decomposition not forced; additive form unjustified; ΔV_Σ approximation inadequate; dominance regimes trivially correct) are absorbed into the Epistemic Status block, the §Discussion “Computational-extraction-gap framing” sub-section, and the missing-channels Working Note above. Originating spike is absorbed archaeology, not a live reference.

- **Reproducibility:** the simulation-check outcomes stated in the Epistemic Status block were produced by the drifting-multi-armed-bandit driver `spikes/sim-three-way-tradeoff.py` (five configurations). The body states the outcomes self-contained; this pointer is reproducibility metadata, to be released as a citable supplement at publication.
- **Open: the unified vs two-stage question.** Simulation suggests the unified approach may be strictly better. If so, the segment should present the unified formulation as primary and the two-stage decomposition as a computational simplification.
- **Relationship to working memory.** Deliberation requires maintaining internal state. For LLM agents, this manifests as context-window competition: tokens spent “thinking” are tokens unavailable for representing the problem state. This architectural constraint on deliberation is real but domain-specific.

Incidental audit gold (lift 2026-05-31). Cross-audit “wandering thoughts” / §14-ideation harvested from the de-novo auditors’ working dirs, deduplicated across substrates and attributed by substrate + audit. Orthogonal pedagogical / framing / figure / naming material, kept separate from the certified theory-fix findings (handled elsewhere). **Coverage:** 3 dirs carry a dedicated reflection (526815, 773921, 829314) plus 451729 covers it inside a Section-II-close batch (01-batch-13). Most other contributing dirs explicitly skipped the §II remainder (this is the final §II segment). Attribution inferred from voice where not explicit, uncertain cases hedged.

1. Candidate Brief prose / pre-prose.

- The deliberation stopping rule in plain language: “stop thinking when the marginal value of your thoughts drops below the rate at which the world is changing” (Claude, AUDIT-WORKING-773921) — i.e. “you must stop thinking and start acting the exact moment your internal ‘aha!’ rate drops below the world’s drift rate” (Gemini, AUDIT-WORKING-829314). A Feynman-criterion-grade gloss of the first-order condition.
- The three activities as one budget, three extraction targets: **“Exploitation:** burn tempo to extract value from reality. **Exploration:** burn tempo to extract data from reality. **Deliberation:** burn tempo to extract latent data from your own internal model” (Gemini, AUDIT-WORKING-829314). A compact triad worth keeping near the Brief.
- “Deliberation as internal exploration in model-space bounded by the computational extraction gap” — the one-line characterization two substrates flagged as the segment’s best contribution (Claude, AUDIT-WORKING-773921; Codex/Claude, AUDIT-WORKING-526815).

2. Candidate Discussion.

- **Analysis paralysis as the formal pathology of crossing the deliberation threshold** (*strongest framing here*). “An engineer spends 3 weeks designing the perfect microservices architecture (Deliberating). But during those 3 weeks, the product requirements changed ($\rho_{\text{delib}} \cdot \Delta\tau$). The resulting architecture is beautiful but useless. The engineer mathematically violated the deliberation threshold. They should have stopped thinking on day 2 and built a prototype to gather external causal information (CIY)” (Gemini, AUDIT-WORKING-829314). A vivid worked illustration of the FOC’s failure mode the segment could carry directly.
- **Why agents (and minds) must “sleep” / “think.”** “If your model M_t is highly complex (a deep neural network or a human neocortex), you physically cannot extract all the useful inferences from it instantaneously. You possess latent knowledge that you haven’t connected yet. Deliberation is the process of running internal queries (simulations, counterfactuals) against your own model to surface that latent knowledge and compile it into a better strategy Σ_t ” (Gemini, AUDIT-WORKING-829314). A candidate Discussion gloss on the computational-extraction-gap framing already in the body. (*Early-conflation texture: the “must sleep” reach is analogy, not derived.*)
- **The cascade as the control-flow that navigates the three-way allocation.** The two segments compose: `#der-orient-cascade` gives the forced ordering; this segment gives the resource allocation across the steps of that ordering. “The Orient Cascade provides the formal algorithmic control flow that navigates the 3-way resource allocation problem defined here” (Gemini, AUDIT-WORKING-773921). A candidate framing-level connector between the two §II-closing segments.

3. Follow-up items.

- **A strategy-structured simulation environment to actually show deliberation’s value.** The Epistemic Status honestly reports the bandit oracle “almost never chose deliberation” because the bandit lacks strategy structure. Multiple substrates flagged the need for “a simulation environment with a deep strategy DAG (like a maze or a coding task) to empirically prove deliberation’s value” (Gemini, AUDIT-WORKING-829314; Claude, AUDIT-WORKING-451729 batch-13 — the bandit is unfavorable by design; a planning-heavy domain would likely show the opposite, already partly noted in the existing “conditions for deliberation dominance” Working Note). A candidate follow-up spike, complementing the existing spikes/sim-three-way-tradeoff.py.
- **Promote the “deliberation as internal exploration” reframe out of Discussion.** Two substrates independently judged that the segment’s best insight is buried: “the ‘deliberation as internal exploration’ reframing is the most important contribution and should be in the Formal Expression or Epistemic Status, not buried in Discussion” (Claude, AUDIT-WORKING-451729 batch-13; Gemini, AUDIT-WORKING-829314 — “shouldn’t be buried under citations”). A structural-placement note for the eventual promotion pass.

4. Readers often ask / wonder.

- **The Active Inference comparison reads as defensive bloat to a fresh reader.** A fresh reader experienced the EFE / Active-Inference comparison paragraph as “a massive, defensive comparison” that “severely damages the readability of the core theory” and “buries” the brilliant core insight under citations; the suggested remedy was to move the AI comparison to an appendix and visually highlight the domain-instantiations table instead (Gemini, AUDIT-WORKING-829314). Preserved as a reader-experience signal about placement/proportion (not a correctness finding — the comparison content itself checked out cleanly for other substrates, e.g. Gemini AUDIT-WORKING-773921 called the KL-direction-from-internal-regret handling “masterful”).
- **When does extra inference-time compute actually help?** Readers reach for the practical corollary: giving a model more “thinking time” yields value only if (1) the model possesses latent knowledge relevant to the problem (high M_t capacity) and (2) the problem is relatively stationary (low ρ_{delib}); in high- ρ environments “inference compute is wasted; you must buy physical tempo instead” (Gemini, AUDIT-WORKING-829314). A candidate readers-often-ask answer with direct contemporary resonance (inference-time scaling).

5. Candidate figures.

- **Three channels from one finite cycle budget + the marginal-benefit stopping rule.** A figure drawing exploit / explore / deliberate as three channels drawing on the *same* finite cycle budget, with the deliberation stopping rule shown as a marginal-benefit comparison — highlighting that deliberation’s benefit is *internal and model-bounded* while exploit/explore both spend *external action opportunities* (Codex/Claude, AUDIT-WORKING-526815).
- **Highlight the domain-instantiations table.** Two substrates called the Military / RL / Software / Organization / AI-coding-agent instantiation table the segment’s strongest pedagogy and recommended it be visually foregrounded rather than left as a trailing table (Gemini, AUDIT-WORKING-829314; Gemini, AUDIT-WORKING-773921 — “phenomenal”).

Belongs elsewhere.

- **LLM inference-time scaling (e.g. o1-style “thinking time”)** ($\rightarrow$ 03-llm-core/ *deliberation budgeting*). The deliberation-threshold math gives a direct account of when test-time compute pays off vs. is wasted (the high- ρ “buy physical tempo instead” conclusion above); the logogenic-substrate chapter is the natural home for the LLM-specific instantiation (Gemini, AUDIT-WORKING-829314).
- **A single framework unifying RL (exploit/explore) with classical planning (deliberation)** ($\rightarrow$ *framing-level / positioning*). “It provides a single mathematical framework that unifies Reinforcement Learning (Exploit/Explore) with Classical Planning (Deliberation)” (Gemini, AUDIT-WORKING-773921). A positioning claim for README / Part-II framing prose, not this segment’s body.

Off-ramp (NOT gold) — routed for adjudication. Certified-track (structure), not pedagogical gold. Recorded so they are not lost; they need strengthen-first / consolidated-queue routing.

- **The first-order condition appears to omit a term.** The objective is $\Delta\eta^*(\Delta\tau) \cdot \|\delta_{\text{post}}\| + \Delta V_{\Sigma}(\Delta\tau) - \rho_{\text{delib}}\Delta\tau$, so its $\Delta\tau$ -derivative carries $\Delta\eta^* \cdot \frac{d\|\delta_{\text{post}}\|}{d\Delta\tau}$ *unless* $\|\delta_{\text{post}}\|$ is held fixed by definition. The stated FOC drops this term; the segment should state which quantities vary with deliberation time and which are fixed baselines (F85 candidate, Codex/Claude AUDIT-WORKING-526815). This touches the segment's one genuinely formal contribution.
 - **The control-regret ceiling should also caveat policy-class expansion.** The body caveats that the δ_{regret} ceiling applies only to the strategy-revision channel and names the objective-revision and uncertainty-reduction escapes; it should *also* name that deliberation which expands Π , changes representation, discovers a new action abstraction, or escalates the continuation convention can move the ceiling (F86 candidate, Codex/Claude AUDIT-WORKING-526815). A scope completion of the ceiling claim.
 - **Symbol overload on λ .** “As $\lambda \rightarrow 0$ ” in the boundary conditions is ambiguous — λ is used for exploration pricing here and for forgetting in nearby strategy-dynamics segments; the intended quantity should be named explicitly (F87 soft candidate, Codex/Claude AUDIT-WORKING-526815).
-

10.3 Discussion: Additional Implications & Discussion

The chapter closes Part II by tying together what the cascade does, what forces it to take the shape it does, and what the cascade looks like at two structural extremes — when the agent is *too confident* in a drifting world (where the cascade has a Lyapunov-forced exploration drive its information-value cousins cannot supply) and when the agent is *coupled-architecture* by construction (where the cascade cannot be produced in cascade order by the model itself and must be recovered at the loop level). The chapter is small in segment count but its results connect Part II to Part III via the directed-separation architectural classification, to 03-11m-core/ via the Class-3-Coupled scaffolding requirement, and back to Part I Ch.4 via the bandwidth floor that deliberation does not escape.

The Orient Cascade




Figure 5: The orient cascade: a five-step resolution order forced by information dependency, not design. The 2x2 diagnostic at step 3 routes the agent to revise the strategy, to escalate before revising the objective, or to terminate at success or capacity limit; a revised model re-enters at step 1.

The cascade is forced by information dependency, not designed.

#der-orient-cascade carries the chapter's central architectural move: the Orient phase of OODA expands into a multi-step cascade — reduce $\delta_{\text{epistemic}} \rightarrow$ assess attainability $\rightarrow$ evaluate strategy $\rightarrow$ (if needed) revise objective — whose *ordering is not a design choice but a consequence of which quantities require which others*. Each step's input depends on the prior step's output. M_t enters every subsequent formula, so it must be updated first; attainability A_O is computed against M_t , so it precedes strategy evaluation; strategy revision needs A_O to know whether the problem is the strategy or the goal; objective revision is last because it requires both the model and the strategy assessment to know that no policy in Π can deliver the objective. The cascade is *derived* from information dependency, not engineered — and that is the framework's load-bearing position against alternative cascade orderings (objective-first, strategy-first, or all-at-once) that have appeared in adjacent literatures.

The downstream consequence is that Part II's diagnostic core — the satisfaction-gap and control-regret split from Ch.3 — has its corrective machinery here. The 2x2 cell map routes to four different revision targets via the cascade: M_t first under epistemic mismatch; Σ_t under positive control regret; Π expansion under control regret persisting after strategy revision; and O_t only when satisfaction gap persists after all earlier revisions. Objective revision is the *last resort*, not the first response to unmet goals. This is the chapter's implications-side payoff: the diagnostic from Ch.3 is what tells the agent *what* to revise; the cascade from this chapter is what tells the agent *in what order*; and the orient cascade's forced ordering is what makes the diagnostic actionable rather than merely informative. Part II would have been a collection of diagnostics without revision machinery without it.

The cascade also depends on the directed-separation scope condition from `#der-directed-separation` (Part II Ch.1) for the *sequential resolution* to be meaningful. For any **causally-disciplined** agent — one running at $\kappa_{\text{processing}} = 0$, whether by construction (Class 1) or in effect (an idealized Class 2 at $\kappa = 0$) — the cascade runs sequentially because the belief-update has no causal dependence on G_t , so the resolution order is a sequential one rather than a simultaneous fixed-point; this consumes only the zero, not the architecture-inspection certificate. As $\kappa_{\text{processing}}$ leaves zero the resolution acquires a coupled component, and in a Class 3 (Coupled) architecture the cascade lives in one merged computation and the resolution order cannot be produced by the model itself. The framework’s prescription for Class 3 — that scaffolded loops around coupled components recover the cascade ordering at the wrapper level — depends on this chapter’s derivation: what scaffolding *is* structurally doing is reproducing the orient cascade externally when the component cannot reproduce it internally. The implication is concrete and significant for practitioners: agentic systems (loops around LLMs) are not engineering convenience; they are the structural mechanism by which Part II’s persistence guarantees become available to Class 3 (Coupled) substrates.

Exploration at two ends of the uncertainty spectrum.

The chapter’s appendix-side derivations carry one of AAT’s load-bearing methodological positions: exploration is forced by Lyapunov persistence collapse, not by epistemic value-of-information. `#deriv-causal-ib-exploration`’s Findings cover this in full — the maximum tolerable observation-noise expectation $U_o^{\max} = U_M(Rc_{\min}/\rho^{\text{eff}} - 1)$ scales with U_M , so as the agent becomes confident ($U_M \rightarrow 0$), the tolerable noise tightens, and the agent is forced to actively seek pristine, low-noise observations to penetrate its own stubbornness. The framework therefore carries *two parallel exploration drives at opposite ends of the uncertainty spectrum*: the standard epistemic drive $\lambda_{\text{info}} \propto U_M$ (explore when uncertain — dominant when U_M is high) and a Lyapunov-survival drive $\lambda_{\text{surv}} \propto 1/U_M$ (explore when confident in a drifting world, because the agent’s own update gain has dropped below the rate at which reality is shifting — dominant when U_M is low). The implications-side payoff that the segment’s Findings flags but is worth repeating is that this *sidesteps active-inference’s dark-room objection at structural rather than parametric level*: an EFE agent with strong priors can stop exploring, an AAT agent in a drifting world physically cannot, because exploration is forced by the same Lyapunov machinery that grounds the rest of Part II.

`#deriv-causal-ib-lmi` carries the matrix-form sharpening. The scalar form of the survival imperative admits a “blank wall” attack — an action that selects a low-noise channel orthogonal to the drift direction satisfies the scalar constraint while the agent’s mismatch in the drifting subspace grows unbounded. The LMI lift to a Linear Matrix Inequality on the Fisher Information Matrix, with a positive-semidefinite matrix Lagrange multiplier $\Lambda \geq 0$ as directional shadow price, resolves the attack by complementary slackness on direction — the exploration bonus for blank-wall actions is mathematically zeroed. The result also makes precise that the Fisher Information Matrix is doing *three structurally distinct* loads in the framework: credit assignment via `#deriv-fisher-whitened-update-rule`, coordinate-forcing via `#disc-additive-coordinate-forcing`’s fourth instance, and now survival via the LMI here. The same geometric object recurs because the framework’s coordinates are forced rather than chosen; the LMI form is the first explicit cross-instance composition where two of those loads — coordinate-forcing and survival — interact within the same machinery.

The composition with Part I Ch.4’s persistence-information-rate floor is worth surfacing. The persistence-bandwidth-floor $C \geq \mathcal{T}/2$ is what a confident agent in a drifting world is structurally required to *meet*; the survival-imperative LMI is what tells the agent *in what directions* to spend the bandwidth. Without the LMI, the bandwidth could in principle be supplied in irrelevant directions; with the LMI, complementary slackness routes the bandwidth toward the drifting modes. The two results compose into a directional version of Part I Ch.4’s central inequality: *bandwidth-per-direction is the binding constraint*, and the LMI is what makes the directionality structurally enforceable rather than implicit.

Coupled architecture, the bias bound, and the scaffolding requirement.

The chapter is also where the framework’s treatment of Class 3 (Coupled) agents lands its most consequential structural commitment. #der-directed-separation’s Findings section covers the architectural classification in full — Class 1 (Separated), Class 2 (Partial), Class 3 (Coupled) — with the explicit Class 3 scope exit that hands the agents off to 03-llm-core/ rather than treating directed separation as an unenforced approximation. #deriv-observation-ambiguity-bias-bound’s Findings cover the universal constant under the (PI) parameterization-invariance axiom: Track 1 (transport inequality, linear in I) gives $C_{W_2}^2 = 2L_{\text{post}}^2 / \rho_{\text{LSI}}$ under log-Sobolev + Lipschitz-posterior; Track 2 (Fisher-Rao, $\sqrt{I}$ scaling) gives universal dimension-free $C_{FR} = \sqrt{2}$ under (PI)+Čencov + small-information regime. The no-go counterexample (heteroscedastic-Gaussian, demonstrating that no universal C exists under Euclidean-parameter norm) is what forces (PI) from convergent representational choice to load-bearing for theorem-level status — adding the fourth instance to #disc-additive-coordinate-forcing’s catalog of forced coordinates (the (PI)+Čencov downstream-theorem cluster). It is *not* a fourth identifiability-floor instance, however — its external theorem (Čencov 1982) *forces* the escape’s uniqueness rather than *forbidding* the inferential task, and the single forced escape (adopt (PI)) is logically incompatible with the ≥ 2 -distinct-escapes structure that defines a floor; the genuine fourth #disc-identifiability-floor instance is #der-architecture-noidentifiability (architecture-noidentifiability from on-policy summary data, dual-anchored on Kalman-Ho 1966 / CHT-at-agent-as-SCM, integrated 2026-05-21).

What the implications segment can develop beyond the source-segment Findings is the *coupled-diagnostic framework* that the bias bound enables, and which does not currently have its own segment home. For Class 2 (Partial) and Class 3 (Coupled) agents — LLMs, transformer-based agents with attention processing goals and observations together — the orient cascade’s information-dependency-forced ordering cannot be produced in cascade order by the model itself; the cascade lives in one merged computation. But the Part II diagnostic quantities ($\delta_{\text{sat}}, \delta_{\text{regret}}$) remain *defined* on post-update state, and the Lipschitz bound $|\delta_{\text{sat}}^{(\text{coupled})} - \delta_{\text{sat}}^{(\text{clean})}| \leq L_A \cdot \|\Delta M_{\text{bias}}\|$ shows the diagnostic error is bounded by the bias bound, which is itself bounded by $\kappa \cdot \mathcal{A}$ (coupling $\times$ ambiguity). The structural conclusion is the load-bearing implication: **agentic systems — scaffolded loops around coupled components — are not just engineering convenience but a structural requirement to recover Part II’s persistence guarantees for Class 3 (Coupled) substrates.** The cascade ordering happens at the loop level when the model cannot produce it; the scaffolding is doing real structural work rather than aesthetic work. For a production LLM-agent practitioner running structured-reasoning templates, multi-step prompts, external monitors, or retrieval-augmented calls, the framework’s claim is that what is being built is *necessary* (not contingent) for AAT-grade behavior in a Class 2 / Class 3 substrate. The cross-domain transfer to human bicameral / dual-process models (Kahneman 2011) and to organizations where executive scaffolding recovers cascade ordering across coupled individuals follows the same structural mechanism.

The engineering anchor worth carrying is the $\kappa \cdot \mathcal{A}$ decomposition. The architectural coupling κ is *fixed* for a Class 3 substrate (you cannot un-couple an LLM); the ambiguity $\mathcal{A}$ is *variable* (you can engineer the prompt and environment). Prompt engineering, in the framework’s vocabulary, is literally $\mathcal{A}$ -reduction — and the bias bound says reducing $\mathcal{A}$ tightens the bound proportionally, with the constant C now bounded by Track 1 or Track 2. The framework predicts (and the empirical pattern supports) that LLMs excel at low- $\mathcal{A}$ tasks (coding, structured reasoning, well-defined transformations) and struggle with high- $\mathcal{A}$ tasks (open interpretation, ambiguous goals, contested values). The framework’s prescription is concrete: where κ is structurally large, design $\mathcal{A}$ structurally small.

Deliberation is internal exploration that does not relax the bandwidth floor.

#disc-exploit-explore-deliberate extends the standard exploit / explore tradeoff into a three-way framing by surfacing deliberation as *internal exploration in model space rather than environment space*. The implications-side payoff is structural: deliberation is Pearl-*do* on a simulated trajectory rather than on the agent’s singular causal trajectory, and the chapter is where this characterization first becomes formally well-posed — Part II Ch.2’s loop-Level-2 result is what makes the deliberation-as-internal-exploration framing distinct from “the agent thinks more” rather than a verbal re-framing.

The composition with Part I Ch.4's bandwidth floor is the finding that has no other natural home: *deliberation does not relax the persistence-information-rate floor; it only changes channel allocation*. The agent that deliberates is internally rerouting its information-processing budget from environmental observation channels to internal-simulation channels, but the total Shannon information rate required to maintain bounded mismatch ($\dot{R} \geq n\alpha/2$ nats/time per dimension) is unchanged. An agent that deliberates indefinitely is *not* exempt from receiving observations at the floor rate — environmental drift continues during deliberation, and the post-deliberation mismatch grows at rate $\rho_{\text{delib}} \cdot \Delta\tau$ per #der-deliberation-cost's drift term. The extended deliberation threshold from #disc-exploit-explore-deliberate formalizes this: the agent should deliberate until the marginal benefit (epistemic gain + strategic gain) equals the marginal drift cost. The framework's posture is that deliberation is *temporary reallocation*, not substitute for external bandwidth — and the cross-domain transfer follows naturally to chess engines (search depth and observation freshness must both be funded), military command (war-gaming and reconnaissance must both be funded), and science (theory-building and experiment-running must both be funded).

The connection to Part II Ch.4's strategy-complexity-cost (9.9) is also worth surfacing here. Part of the cost of deliberation is the cognitive cost of holding Σ_t in the agent's representational capacity while operating on it. For an LLM agent, this means fitting the strategy in the context window — and the framework's claim is that this is the same constraint that bounds online operation, not a separate deliberation-mode constraint. Deliberation buys epistemic and strategic improvement at the cost of cycle budget *and* representational capacity simultaneously; both are bandwidth-flavored constraints, and the persistence-bandwidth-floor framework subsumes them. This is one of the chapter's quieter findings — the framework's bandwidth vocabulary unifies environmental-observation bandwidth, strategy-representation capacity, and deliberation budget under a single Shannon-rate floor.

What Part II closes on, and the bridge into Part III.

The chapter is the structural close of Part II, and what Part II has delivered as a whole is now worth naming. The diagnostic core (satisfaction-gap and control-regret split, convention-indexed by the C1/C2/C3 hierarchy) gives the framework two orthogonal signals that distinguish “the world doesn't permit it” from “you're not doing it well enough.” The strategy-DAG commitment with Markov factorization under causal sufficiency gives the formal representation the diagnostic operates on. The Pearl-Level-2 access from the loop, combined with the Correlation Hierarchy's scope statement and the chapter's directed-separation architectural classification, gives the agent the data substrate and the architectural class it sits in. The bias bound for Class 3 (Coupled) agents, with its conditional-theorem constant under named sub-scopes, gives the framework's quantitative reach into LLM-class substrates without overclaim. The orient cascade gives the corrective machinery; the survival-imperative drive gives the structural exploration mandate; the coupled-diagnostic framework gives the scaffolding-as-structural-requirement claim.

The result of all this is what #result-section-ii-survival formalizes: Part II's definitional and structural architecture transfers exactly to fully-coupled (Class 3) agents in 16 of 24 cases, with the remaining results either approximating with bounded error (5/24) or naming the precise modification required (2/24 plus 1 boundary-defining failure). The bias-bound theorem under named sub-scopes is what lifts the approximate-survival category from heuristic to conditional theorem. Part II's reach is wider than the most-restrictive scope condition suggests; this chapter is where the chapter-end can name that reach without overclaiming.

The bridge to Part III is direct. The directed-separation classification opens to *axis-decomposed composite-level class inheritance* (per the table in #der-directed-separation §“Composite-level class inheritance”): composite architectural class lifts through routing structure plus substrate sharing, *not* through goal alignment per se. Goal alignment vs strategic interaction operates on the parallel-but-independent *dynamic-regime* axis (#disc-dynamic-regime-axis), where strategic composition transitions composites from R0 contraction-regime to R1 equilibrium-regime or R2 cyclic-distributional-regime without architectural-class change (under goal-blind routing). The third instance of the identifiability-floor pattern (Instance 3, broadened 2026-05-21 to the coupling-topology bit) lives in Part III: composite contraction *and regime* certification from component-marginal data is structurally impossible (Liberzon 2003 common-Lyapunov-nonexistence;

cf. `#deriv-regime-marginal-indistinguishability` for the broader regime-axis witness), with five structural escapes including observer-on-sub-agent interventional access (Mode 2 of the loop-Level-2 pattern from Ch.2) and the new passive escape (e) of composite-level convergence-rate-class observation. Part III is where AAT's machinery operates at the composite scope, with the chapter-end here as the last single-agent prerequisite to land.

What this chapter closes on, then, is the structural completion of Part II: the cascade that makes the diagnostics actionable, the dual-drive exploration that prevents dark-room collapse, the bias bound that quantifies Class 3 (Coupled) reach with conditional-theorem precision, the scaffolding requirement that makes agentic systems structurally non-optional rather than aesthetic, and the deliberation framing that grounds the three-way resource allocation in the same Shannon-rate floor as everything else. Part III's machinery operates on what Part II built; the implications work here is the last single-agent-scope synthesis before the composite chapters take over.

Working Notes.

- **Causal-discipline re-scope (2026-05-31).** The sequential-resolution clause was re-scoped from “Class 1 (Separated)” to *causally-disciplined* (the $\kappa_{\text{processing}} = 0$ object of `#der-directed-separation`): the cascade runs sequentially rather than as a simultaneous fixed-point because f_M has no G_t -dependence — that consumes only the zero, held equally by Class 1 by construction and an idealized Class 2 in effect, not the architecture certificate. Left untouched: the Class-3 scaffolding clause (cascade recovered at the loop level — correct and class-specific), and the cascade *ordering*-exactness in `#der-orient-cascade`, which is forced by information dependency and is not κ -scoped at all (no class label appears on it, correctly — so `#der-orient-cascade` needed no edit). The $\varepsilon(\kappa)$ for the sequential resolution is *inherited, not new*: off $\kappa = 0$ the cascade's step-1 output M_{t+} is displaced by the $\Delta M_{\text{bias}} \leq C\kappa I$ bound of `#deriv-observation-ambiguity-bias-bound`, and every downstream step reads that displaced M_t ; the resolution becomes one iteration of a coupled fixed-point whose residual is governed by that same bound. A dedicated propagated-residual bound (how the per-step bias compounds across the five-step cascade) is an open strengthening, not attempted here.
- This segment is a chapter-end discussion grouping findings from `#der-orient-cascade`, `#disc-exploit-explore-deliberate`, `#der-directed-separation` (the cascade depends on it for the sequential-resolution scope), `#deriv-observation-ambiguity-bias-bound` (Appendix A, the bias-bound constant), `#deriv-causal-ib-exploration` and `#deriv-causal-ib-lmi` (Appendix A, the survival-imperative drive and its matrix lift), and `#der-deliberation-cost` (Part I Ch.4, the deliberation threshold this chapter extends). The closing-Part-II reading in §5 leans on `#result-section-ii-survival` which is the formal classification of which Part II results transfer to Class 3 (Coupled) agents at exact / approximate / modified status.
- **Cross-reference vs canonical-home policy.** Four source segments in this chapter's scope carry well-written `## Findings` sections — `#der-directed-separation` (architectural classification), `#deriv-observation-ambiguity-bias-bound` (universal constant under (PI)), `#deriv-causal-ib-exploration` (survival-imperative drive), and `#deriv-causal-ib-lmi` (matrix-form lift). All four are cross-referenced rather than restated. The findings whose canonical home is this implications segment are: `#20` (Orient Cascade as derived from information dependency), `#13` (Coupled Diagnostic Framework — scaffolded loops recover the cascade for Class 3 substrates; no segment currently carries this finding directly), and `#35` (Internal Deliberation Pays Bandwidth Floor — compositional finding with no segment home).
- **#13 (Coupled Diagnostic Framework) — segment exists in 03-llm-core/.** Corrected from earlier draft: `#result-coupled-diagnostic-framework` exists as a segment in `03-llm-core/src/` (where the Class 3 (Coupled) treatment lives), not in `01-aat-core/`. The segment does not currently carry a `## Findings` section, so the implications-side surfacing here covers what AAT Ch.5 unlocks (cascade recovery via scaffolding at the loop level, bias bound via $\kappa \times \mathcal{A}$); the Logogenic chapter-end implications carry the finding from the logogenic-substrate-specific side. A future cycle may add a `## Findings` section to the existing segment in `03-llm-core/`; until then, the chapter-end implications carry the catalog-grade content.
- **Connection to NeurIPS submissions.** The bias-bound result is NeurIPS Paper 3 (`~/src/neurips/03-llm-hallucinate-bound/`) in formal theorem form; the (PI) axiom from `#scope-agent-identity` is load-bearing for the universal-constant route under the no-go in §4 of Paper 3. The cascade-recovery-via-scaffolding argument is a structural claim AAT has been carrying since the Part II reorganization; whether it lifts to theorem grade in any current submission is not flagged. The survival-imperative drive composed with the LMI matrix lift is the formal apparatus behind NeurIPS Paper 1 (“Tragedy of the Confident Agent”, `~/src/neurips/01-tragedy-confident-agent/`).

- **Identifiability-floor pattern coverage update.** With this chapter's content three formal instances are in view: Instance 1 (on-policy L0-insufficiency, Part II Ch.4), Instance 2 (L1' mixture under unobservable common cause, Part II Ch.3), and Instance 4 (architecture-noidentifiability from on-policy summary data, this chapter / Appendix A — #der-architecture-noidentifiability's Kalman-Ho similarity-orbit no-go, dual-anchored with CHT-at-agent-as-SCM). Instance 3 (composite contraction certification from component-marginals via Liberzon common-Lyapunov-nonexistence) lives in Part III. The heteroscedastic-Gaussian no-go in #deriv-observation-ambiguity-bias-bound §4 is *not* a floor instance — it is correctly classified as a downstream theorem of #disc-additive-coordinate-forcing's (PI) commitment (single forced escape via Čencov-uniqueness, logically incompatible with the ≥ 2 -distinct-escapes structure required of a floor). Instance 5 (Mehra non-identifiability in adaptive Kalman) remains candidate; the chapter-end pattern coverage is now five-sixths complete on the planned AAT chapter-end work; the Instance 3 surfacing in Part III Ch.5 completes it.

Incidental audit gold (lift 2026-05-31). Cross-audit “wandering thoughts” / §14-ideation harvested from the de-novo auditors' working dirs, deduplicated across substrates and attributed by substrate + audit. Orthogonal pedagogical / framing / figure / naming material, kept separate from the certified theory-fix findings (handled elsewhere). **Coverage:** 2 dirs carry a dedicated reflection (526815, 773921). Most other contributing dirs stopped before this Part-II-closing chapter-end segment. Attribution inferred from voice where not explicit, uncertain cases hedged.

1. Candidate Brief prose / pre-prose.

- “You cannot ‘think’ your way out of needing sensory bandwidth” — a one-line gloss of the chapter's deliberation-does-not-relax-the-floor result ($\dot{R} \geq n\alpha/2$ holds while the agent deliberates because the environment keeps drifting) (Gemini, AUDIT-WORKING-773921). Compact, vivid, Feynman-criterion-grade.
- “Prompt Engineering is literally Ambiguity Reduction ($\mathcal{A}$ -reduction)” — the $\kappa \cdot \mathcal{A}$ decomposition stated as a practitioner takeaway: κ (architectural coupling) is fixed for a Class-3 substrate, $\mathcal{A}$ (ambiguity) is the engineerable lever (Gemini, AUDIT-WORKING-773921). Already in the body; flagged as the line worth foregrounding.
- Scaffolding “is not a ‘prompt engineering hack’; it is the structural mechanism required to force a Class 3 agent to execute the Orient Cascade sequentially, thereby recovering the Lyapunov persistence guarantees of Part I” (Gemini, AUDIT-WORKING-773921). A candidate Brief framing of the chapter's load-bearing scaffolding claim.

2. Candidate Discussion.

- **Agentic workflows are mathematically superior to raw calls — and why.** Two substrates converged on the chapter's payoff as the strongest engineering takeaway of Part II: scaffolded loops (ReAct loops, external memory, monitor agents, retrieval) are *structurally necessary*, not convenient, because a Class-3 substrate cannot run the cascade in cascade-order internally (beliefs and goals are mixed in one merged computation), so the ordering must be recovered at the loop level (Gemini, AUDIT-WORKING-773921; the bounded form is sharpened by Codex/Claude, AUDIT-WORKING-526815 — see off-ramp). A candidate Discussion sharpening that names *why* the loop recovers what the component cannot.

3. Follow-up items.

- None surfaced beyond the certified-track scope concerns (see off-ramp). The chapter-end's own Working Notes already track the #13 (Coupled Diagnostic Framework) canonical-home question and the #result-section-ii-survival dependency.

4. Readers often ask / wonder.

- **Does the chapter's summary voice outrun its proofhomes?** Fresh readers experienced several proof-bearing claims (survival-imperative exploration, the causal-IB LMI structure, the bias-bound constants, the Section-II survival counts, the scaffolding requirement) as *presented as settled by the chapter* when the proof credit lives in deferred appendix/source homes (Codex/Claude, AUDIT-WORKING-526815). This is an audit-experience signal about a chapter-end implications segment: the cross-reference-vs-restate policy (already in this segment's Working Notes) is doing the right thing, but a reader can still read the synthesis voice as proof. A candidate one-line orientation at the chapter open (“results below are surfaced here; their proofs live at the cited homes”) would preempt the stumble.

5. Candidate figures.

- **Three-zone proof-flow diagram.** A chapter map drawn as three zones — *local anchors already established*, *deferred proof homes* (appendix/source segments), and *architectural prescriptions* — routing the strongest claims visibly *through* the deferred homes rather than appearing to originate in the local synthesis (Codex/Claude, AUDIT-WORKING-526815). Directly addresses the “summary voice outruns proof homes” reader stumble above by making the proof-provenance visual.

Belongs elsewhere.

- **Architecture-noidentifiability as a wall for mechanistic interpretability** ($\rightarrow$ #der-architecture-noidentifiability / *Appendix A, and positioning prose*). The Instance-4 floor (Kalman-Ho similarity-orbit non-uniqueness) “formally proves that you cannot reverse-engineer the true causal architecture of a black-box model just by watching its inputs and outputs, because infinite internal similarity transformations yield the exact same behavior — a massive mathematical wall for Mechanistic Interpretability” (Gemini, AUDIT-WORKING-773921). The canonical home is #der-architecture-noidentifiability; the interp-field framing is positioning reach worth carrying there, not here.

Off-ramp (NOT gold) — routed for adjudication. Certified-track (scope/status — the chapter-end synthesis voice vs. proof provenance), not pedagogical gold. Recorded so they are not lost; several overlap with this segment’s existing cross-reference-vs-canonical-home Working Notes and may already be dispositioned — flagged for the consolidated queue to confirm.

- **Proof-credit should stay with declared homes.** impl-orient-cascade is discussion-grade but states survival-imperative exploration, the causal-IB LMI structure, bias-bound constants, the Section-II survival counts, and the scaffolding requirement as if settled by the chapter (F88 candidate, Codex/Claude AUDIT-WORKING-526815).
 - **The scaffolding-necessity claim is too broad if read literally.** The local argument supports a narrower claim: scaffolding is needed to *recover AAT-grade cascade guarantees under coupled representation/attention/memory/action channels*, not an unconditional “scaffolded loops are structurally necessary for Class-3 substrates” (F89 candidate, Codex/Claude AUDIT-WORKING-526815).
 - **“Deliberation is Pearl-do on a simulated trajectory” is too strong.** Pearl’s *do* is an intervention in an SCM; model-internal simulated deliberation is better described as model-internal counterfactual intervention / policy-model simulation (F90 candidate, Codex/Claude AUDIT-WORKING-526815).
 - **The “deliberation leaves total Shannon rate unchanged” claim needs scope.** Internal computation cannot replace external information indefinitely, but better policies, sensing, compression, and model selection *can* change the effective disturbance, allocation, or required rate (F91 candidate, Codex/Claude AUDIT-WORKING-526815).
-

Part III

Agentic Composites

Chapter 11

Meta-Architecture II: Composition Consistency and the Dynamics-on-Architectural-Class Facet

Part II opened with the certificate spine and its three Part-II-operative facets (M1 boundary / M2 scope-of-existence / M3 forced-identity). Part III opens with two cross-level meta-recognitions that organize the composite-agent material that follows: the composition-consistency commitment (how the framework *applies across levels*) and the fourth certificate-spine facet (how architectural class *changes* under operations). Both are introduced before any Part III content chapter because both are framework-level recognitions — the composite-agent machinery in subsequent chapters is the instantiation, not the source.

Composition consistency — cross-level compatibility as framework commitment (11, methodological commitment). The framework's scope conditions (1.5, 1.6) do not restrict which level of description AAT applies to; the same observable can be analyzed at the individual / team / organizational level, and the theory must not return contradictory answers across those analyses. The commitment has three operational laws (tempo composition, persistence compatibility, mismatch consistency) and a three-layer decomposition (scope selection — 12.2; admissibility — 13.1; transfer — Tier-dependent via AI). Under Tier 1M sub-agents the composite contraction rate admits a closed form; outside Tier 1M the screening test $\tau_{eq} \ll \tau_{ext}$ is heuristic. Cross-level companion to the four certificate-spine facets — the spine names the geometry of a single certificate, composition-consistency names the relationship across certificates at different levels. Sister to 11.1: M4 names how architectural class *changes* under operations; composition-consistency names how the framework *applies across levels* regardless of architectural class.

M4 — Modularity-state dynamics (11.1, dynamics-on-architectural-class facet): modularity is a *contested property* — composite agents do not sit fixed at one architectural class; they move under three structurally-distinct operations. **Truthification** (self-driven, modularity-increasing; restorative valence): the wrapping construction at the component level (13.4, in Ch.2 of this Part) and institutional scaffolding at the composite level (per 11.3 §"Defensive scaffolding"). **Strategic self-coupling** (self-driven, modularity-decreasing; offensive-enabling valence): operation leg 11.2. **Adversarial coupling pressure** (externally-driven, modularity-decreasing; vulnerability-creating valence): operation leg 11.3. Three pairwise dual relationships organize the cluster — truthification $\leftrightarrow$ adversarial-pressure as direct dual (cultivation vs attack); truthification $\leftrightarrow$ strategic-self-coupling as goal-belief-axis dual (both wrapping-construction-style commitments on opposite poles — goal-blind belief-update vs goal-conditioned action-selection); strategic-self-coupling $\leftrightarrow$ adversarial-pressure as same-architecture-shape-opposite-driver. M2's Architecture ladder reads as the phase space across which M4 operations move populations.

The two operation-leg segments. 11.2 develops the self-driven coupling-increasing operation with three structural extensions ((M1) coupling-dependent action space, (M2) strategy-DAG enabling-edge type, (M3) reversibility-cost asymmetry) and a four-mechanism prior-art adoption (Schelling 1960 external commitment devices, Ainslie 1992/2001 intertemporal personal rules, Akerlof-Kranton 2000/2010 identity coupling, Frank 1988 emotional commitment). 11.3 develops the externally-driven coupling-increasing operation — adversaries drive opponent coupling because coupling expands attack surface (identity-binding / affect / sunk-cost mechanisms); diagnostic capture is directional; orient-cascade inversion under pre-loaded O_t ; defensive scaffolding as composite-agent restoration. The two operation legs are opposite-valence operations on the same architectural-class state $\kappa_{\text{processing}}$ per 6.3.

This is a *meta-recognition* about how AAT applies, not a fact about how any individual agent works: AAT's predictions must be compatible across levels of description. If a system S and any decomposition $\{S_1, \dots, S_n\}$ both satisfy the scope condition, the theory's claims at the S -level cannot contradict its claims at the $\{S_i\}$ -level. The scope condition does not restrict which level the theory applies to, so the theory itself must be level-invariant — otherwise the same observable could receive contradictory answers depending on where the analyst draws the modeling boundary. The commitment is methodological at the meta-requirement layer (a framework choice about applicability, not a physical fact about agents); its operational consequences are tiered.

This commitment is the cross-level companion to the four certificate-spine facets (5). Where the spine's facets describe the geometry of a *single* equilibrium-stability certificate (interior / scope-of-existence / forced-identity / boundary / projection), composition consistency describes the relationship *across* certificates at different levels — the same persistence argument must apply at every level where the scope condition is met, and the operational unpacking below is how that “applies at every level” claim cashes out. It sits alongside 11.1 as the second cross-level meta-recognition: M4 names how architectural class *changes* under operations; composition consistency names how the framework *applies across levels* regardless of architectural class. Together they organize the Part III meta-architecture.

The commitment has formal teeth: three composition laws must exist. *Tempo* at the composite level must be expressible as a function of sub-agent tempos plus their coordination structure; *persistence* of the composite must be derivable from the sub-agents' individual persistence plus that coordination structure; *mismatch* must compose consistently. The laws themselves — what tempo and persistence look like at the composite level, and when they transfer cleanly — are *derived* from this postulate in Part III under specific conditions, organized into three successively more specific layers:

- a *scope-selection* layer (12.2, teleological-alignment threshold) — *which* decompositions qualify as composites at all;
- an *admissibility* layer (13.1, conditions (A1)–(A4) and the bridge lemma) — *when* composite macro-dynamics are well-posed under those scope conditions;
- a *transfer* layer (13.1's tier-specific contraction assumption together with AI's (CC-parallel)/(CC-cascade)/(CC-feedback) closed forms) — *which* Part I/II results lift faithfully, with sharpness depending on the tier of the sub-agents.

The transfer layer is where the deepest result lives. For **Tier 1** sub-agents (Bayesian updaters on exponential families, linear correctors with positive-definite gain, gradient descent on strongly convex losses), the composite contraction rate admits a closed-form lower bound — the parallel case yields $\lambda_c = \min_i \lambda_i$ (the composite is bounded below by its slowest sub-agent), cascade composition gives the same up to a coupling-gain adjustment, and heterogeneous negative-feedback composition gives a closed-form inequality involving feedback-loop gains and coordination costs. Macro-level persistence then reduces to the same persistence inequality used at the individual level, applied to the composite. For Tier 2 (locally convex, nonlinear) and Tier 3 (non-convex, discontinuous) sub-agents, this transfer holds only locally with degradation or per-domain, and the qualitative screening test $\tau_{\text{eq}} \ll \tau_{\text{ext}}$ — internal equilibration timescale short relative to external-dynamics timescale — remains useful but does not entail the formal persistence condition without further verification.

The boundary an analyst draws is itself a modeling choice. A development team is simultaneously an individual developer (an AAT agent), a team (a composite AAT agent), and part of an organization (a sub-agent in a larger composite); the scope condition is satisfied at every level. Composition consistency is what guarantees the theory does not return contradictory answers about observables regardless of which boundary is drawn, and the layered derivation above is what makes that guarantee precise rather than aspirational.

For any system S satisfying the scope condition (1.6), and any decomposition of S into subsystems $\{S_1, \dots, S_n\}$ where each S_i also satisfies the scope condition, AAT's predictions at the system level must be compatible with its predictions at the subsystem level. Specifically, composition laws must exist such that:

Commitment (composition-consistency).

1. **Tempo composition:** $\mathcal{T}_S$ is expressible as a function of $\{\mathcal{T}_{S_i}\}$ and the coordination structure among them
2. **Persistence compatibility:** the system's persistence is derivable from the sub-agents' individual persistence conditions plus coordination structure
3. **Mismatch consistency:** δ_S is derivable from $\{\delta_{S_i}\}$ and their interaction structure

Cross-level compatibility requires three successively more specific conditions, each with its own segment in Part III:

Structural consequence (derivation hierarchy).

1. **Scope:** The decomposition is a *composite* (not merely a multi-agent system) when 12.2 is satisfied — teleological alignment via at least one of three disjunctive routes (shared objective, hierarchical derivation, or mutual benefit). Without scope-satisfaction via any route, level-compatibility is vacuous (composition quantities are not well-defined, so there is nothing to be consistent about).
2. **Admissibility:** Given a composite, its macro-dynamics are well-posed when the closure admissibility conditions (A1)–(A4) hold (13.1). These require AAT-shaped macro-state, macro-mismatch, macro-tempo, and sector-bounded correction at the composite level.
3. **Transfer of results:** Individual Part I/II results lift to the composite level through the bridge lemma in 13.1, which requires a *tier-specific contraction assumption* beyond (A4). For **Tier 1** agents (all Bayesian updaters on exponential families, linear correctors with positive-definite gain, gradient descent on strongly convex losses), the contraction holds and Part I results transfer exactly. For **Tier 2** agents (locally convex, nonlinear prediction models), transfer holds locally with factor degradation. For **Tier 3** agents (non-convex optimization, discontinuous corrections, non-mismatch-driven state), contraction must be verified per-domain and Part I results do not automatically lift.

Composite contraction rate — closed form under Tier 1M. When sub-agents satisfy the contraction-template preconditions (CT1)–(CT3) of AI (Tier 1M, equivalently Tier 1 of 13.1's bridge lemma under the DA2¹-inc $\equiv$ (CT2)-at- $M = I$ equivalence) and the composition topology falls under one of the closure cases of AI, the composite contraction rate λ_c admits a closed-form lower bound:

Derived (Conditional on Tier 1M + admissible composition topology, from #result-contraction-template (CC-parallel) / (CC-cascade) / (CC-feedback)).

- **Parallel composition** (CC-parallel): $\lambda_c = \min_i \lambda_i$ — the composite contraction rate equals the slowest sub-agent's, with composite metric $M_c = \text{blockdiag}(M_1, M_2)$. Equivalently in timescale form, the composite relaxation time $\tau_c = 1/\lambda_c = \max_i \tau_i$ is bounded *below* by the slowest sub-agent's timescale.
- **Hierarchical / cascade** (CC-cascade): same bound, up to coupling-gain adjustment.

- **Negative-feedback heterogeneous** (CC-feedback) / (CM2-M): $(\lambda_1 - C_1)(\lambda_2 - C_2) > k_{12}k_{21}/4$ with feedback-loop gains k_{ij} and coordination costs C_i from 15.1.

Macro-level persistence then reduces to the standard persistence condition (4.5) applied to the composite: $\alpha_c > \rho_{\text{eff}}/R_c$, with α_c bounded by the topology-specific result above and effective disturbance $\rho_{\text{eff}} = \rho_{\text{ext}} + \varepsilon^* \nu_c$ from 13.2's bridge-lemma instantiation. Both Part I results (single-agent persistence) and the bridge-lemma's macro-tracking faithfulness flow from this single inequality. This is the formal counterpart to the screening test below: under Tier 1M the screening test is *derived*, not heuristic.

Practical screening test for composites outside Tier 1M. Tier 2 sub-agents (extended-Kalman, locally-convex gradients, nonlinear prediction models) and Tier 3 sub-agents (non-convex optimization, discontinuous corrections, non-mismatch-driven state) do not in general carry the closed-form composite-rate bound above — the bridge-lemma contraction holds only locally with $\kappa(D\delta)^2$ degradation (Tier 2) or must be verified per-domain (Tier 3). For these composites, the qualitative condition

Heuristic (timescale separation — residual scope: Tier 2 / Tier 3).

$$\tau_{\text{eq}} \ll \tau_{\text{ext}} \quad (11.1)$$

— internal equilibration timescale (the time for sub-agents to approximately synchronize models and coordinate actions) short relative to external-dynamics timescale — remains a useful *discussion-grade screening test*. When it holds, composite description is plausibly well-posed; when it fails, composite description is plausibly broken. The screening test is the operational analog of the persistence condition (4.5), but for Tier 2 / Tier 3 it does not entail the formal persistence condition without local-region or per-domain verification. Whether common organizational settings (software teams, military units) genuinely sit in Tier 1M or merely satisfy the heuristic without satisfying Tier 1M conditions is an empirical question, not a derived fact.

Axiomatic for the meta-requirement (cross-level compatibility). The postulate itself is a structural requirement for AAT's internal consistency: if the scope condition doesn't restrict which level the theory applies to, the predictions must not contradict across levels.

The operational consequences decompose into three layers with distinct epistemic statuses:

- **Scope selection** (12.2, robust-qualitative): which decompositions constitute composites. The scope condition is a disjunction of three qualitatively distinct routes (shared objective, hierarchical derivation, mutual benefit); whether they reduce to a common scalar threshold is open. Coverage of well-understood cases holds via the disjunction.
- **Admissibility of composite dynamics** (13.1, conditional): the (A1)-(A4) conditions are stated formally; the bridge lemma is conditional on the incremental sector bound (DA2'-inc), strictly stronger than (A4) alone.
- **Transfer of individual-agent results** (tier-dependent): Tier 1 composites transfer Part I/II results exactly; Tier 2 transfer with degraded factors; Tier 3 require per-domain verification. This is the sharpest scoping of “every result applies at every level” — it applies at every level *for Tier 1 composites*, degrades gracefully for Tier 2, and holds per-domain for Tier 3.

The composite contraction rate has two epistemic layers. **Tier 1M is derived (exact)**: under (CT1)–(CT3) on each sub-agent and an admissible composition topology, λ_c admits the closed-form lower bound from A1's (CC-parallel) / (CC-cascade) / (CC-feedback) — the slowest-sub-agent / coupling-adjusted / heterogeneous (CM2-M) cases respectively. The screening condition $\tau_{\text{eq}} \ll \tau_{\text{ext}}$ is, in this regime, the formal persistence condition $\alpha_c > \rho_{\text{eff}}/R_c$ specialized to internal-vs-external rate balance — not heuristic. **Tier 2 and Tier 3 retain heuristic / discussion-grade status**: the closed-form composite rate does not transfer (Tier 2 with local degradation; Tier 3 per-domain), and $\tau_{\text{eq}} \ll \tau_{\text{ext}}$

remains a useful screening test without entailing macro-level persistence. The “small gap between screening and Tier 1M conditions in common settings” is an empirical claim about the population of real composites, not a derived fact.

The same pattern as individual persistence. The persistence condition says: there is a measurable threshold below which an individual agent’s mismatch grows without bound and the agent degrades. Composition consistency says the same thing at the composite level: there is a measurable threshold below which the composite description breaks down. Under Tier 1M, the threshold is exact and closed-form — $\alpha_c > \rho_{\text{eff}}/R_c$ with α_c from (CC-parallel) / (CC-cascade) / (CC-feedback). Under Tier 2 / Tier 3, the threshold is only qualitatively captured by the screening ratio $\tau_{\text{eq}}/\tau_{\text{ext}}$ and exact transfer requires local-region or per-domain verification. The two-tier structure is the precise way “the theory applies broadly” — broadly under Tier 1M, qualitatively-with-conditions outside it.

What this buys the theory. With composition consistency stated early: - Part I/II results lift to the composite level for Tier 1 composites passing the composition scope condition, with degraded-factor lift for Tier 2 and per-domain verification for Tier 3 — the three-layer decomposition above makes this precise rather than leaving “applies at every level” as an unbounded claim - Part III becomes the study of *what happens near and beyond the threshold* — coordination overhead, unity dimensions, symbiogenic transitions, adversarial dynamics — rather than a separate multi-agent theory - The formal test for composition validity (13.1) develops the rigorous version of the timescale condition, with the incremental sector bound identifying which agent classes admit the contraction needed for cross-level transfer; AI lifts that bound into a metric-aware framework where (CC-parallel) / (CC-cascade) / (CC-feedback) yield the topology-indexed closed forms cited above

When composition fails. The persistence condition $\alpha_c > \rho_{\text{eff}}/R_c$ (Tier 1M) — or its qualitative shadow $\tau_{\text{eq}} \ll \tau_{\text{ext}}$ (Tier 2 / Tier 3) — fails when: - Internal coordination slows: C_{coord} from 13.2 rises through coupling overhead, conflict, or bureaucratic process, depressing α_c below the threshold; equivalently, τ_{eq} stretches. - External dynamics accelerate: ρ_{ext} rises (adversary acts faster, market shifts, crisis compresses decision timescales); equivalently, τ_{ext} shortens. - Both simultaneously (the classic organizational failure mode — internal friction increases while external demands intensify).

This is the formal analog of Brooks’s Law: adding people to a late project increases $\varepsilon^* \nu_c$ in ρ_{eff} (13.2’s coordination-overhead bound) and stretches τ_{eq} while ρ_{ext} and τ_{ext} stay fixed (the deadline doesn’t move). Under Tier 1M the mechanism is formal — eventually the persistence inequality flips. Under Tier 2 / Tier 3 the mechanism is qualitative. Whether the specific mechanism (coordination-overhead saturation crossing the persistence threshold) is the dominant cause of Brooks’s Law in practice is an empirical question.

The boundary is a modeling choice. A development team is simultaneously: individual developers (each an AAT agent), the team (a composite AAT agent), and part of an organization (a sub-agent within a larger composite). The scope condition is satisfied at every level. Composition consistency ensures the theory doesn’t give contradictory answers about observable quantities (e.g., whether the team persists) regardless of which boundary is chosen.

What composition consistency does NOT say. It does not specify the form of the composition laws — those are derived in Part III (13.2). It does not say every decomposition is equally useful for analysis. And it does not require perfect internal coordination — only that internal equilibration is fast relative to external dynamics.

- **Strengthening attempt — outcome.** The “macro-timescale bounded below by slowest sub-agent” claim was attacked via AI’s compositional theorems before falling back to heuristic-only framing. Under Tier 1M (sub-agents satisfying (CT1)–(CT3) with metrics M_i , rates λ_i): (CC-parallel) yields $\lambda_c = \min_i \lambda_i$ exactly under blockdiag composite metric — the formal version of “composite is no faster than the slowest sub-agent”; (CC-cascade) yields the same bound up to coupling-gain adjustment; (CC-feedback) / (CM2-M) yields the heterogeneous closed-form inequality. No new math required — the strengthening is achieved by binding the postulate’s heuristic to the existing (CC-*) closed forms via the DA2’-inc $\equiv$ (CT2)-at- $M = I$ equivalence in 13.1. Tier 2 (local with $\kappa(D\delta)^2$ degradation) and Tier 3 (per-domain) retain heuristic / discussion-grade screening; this residual is what the strengthening attempt could not eliminate.

- The timescale separation condition is essentially the singular perturbation argument from 4.9 applied to composition: the fast internal dynamics approximately equilibrate, and the composite's behavior is described by the slow (external) dynamics on the equilibrium manifold. The formal connection should be made explicit when temporal-nesting is reviewed.
- Composition of directed separation: if each sub-agent's f_M is G_t -independent, does the composite's f_M^c remain G_t^c -independent? Hypothesis: goal-blindness composes, BUT coordination routing may break it — if which observations reach the composite depends on the shared objective, the composite's effective observation function is goal-dependent. This is the organizational analog of the LLM scope restriction in 6.3.
- The “atomic agent” question: if every agent is decomposable, where does it bottom out? At agents whose internal dynamics are not usefully described by AAT — below the level where observations, actions, and uncertainty exist, the scope condition fails and the recursion terminates.
- The relationship to holons (Koestler 1967): an AAT agent satisfying composition consistency is a holon — simultaneously a whole (analyzable as a single agent) and a part (decomposable into sub-agents). The term is occasionally useful but carries significant mystical baggage from later appropriations. Use sparingly.
- **Provenance — retype from postulate to discussion-grade meta-recognition (2026-05-28).** Originally landed in Volume 1 Part I Ch.1 as a type: postulate / status: axiomatic segment with the slug post-composition-consistency. The OUTLINE annotation (possibly out of place) recorded the placement question through multiple cycles. Audit 384279 (Claude Opus 4.7, 2026-05-27) recommended hybrid relocation — strip operational unpacking back to Part III, keep postulate proper in Ch.1. Joseph 2026-05-28 reframed: meta-segments add machinery and cognitive weight that's *only realized* when given pedagogical work to do, not buried as footnotes; composition consistency is itself a meta-segment in the making, structurally peer to 11.1 as a cross-level meta-recognition (M4 names how architectural class *changes*; composition consistency names how the framework *applies across levels*). Recognized as Meta-Architecture II member; retyped to discussion / discussion-grade; slug renamed post-composition-consistency → disc-composition-consistency. Operational content (Tier 1M closed-form derived block, screening test, Brooks's-Law instantiation) retained intact — now local rather than forward-referenced, since the segment's home is Part III where its dependencies live. The eq-tag *[Commitment (composition-consistency)]* replaces the prior *[Postulate ...]* to reflect the methodological-commitment character (a framework choice about applicability, not a physical-fact axiom). Strengthening cycle — Findings block + prior-art landscape (Simon's near-decomposability, Koestler's holons, multi-level selection, renormalization-group, actor-systems, the Friston-Heins-Verbelen-Da Costa 2025 scale-free active inference machinery, which has since landed in M's promotion 2026-06-10) + paper-extraction candidacy — queued separately in spikes/PROPOSED.md.

Cross-audit “wandering thoughts” / §14-ideation harvested from the de-novo auditors' working dirs, deduplicated across substrates and attributed by substrate + audit. Orthogonal pedagogical / framing / figure / naming material, kept separate from the certified theory-fix findings (handled elsewhere). **Coverage:** the auditors encountered this segment under its *prior* identity post-composition-consistency (Part I Ch.1, type: postulate) before the 2026-05-28 retype-and-relocate to Meta-Architecture II — so the rich finding-stream they produced is the *F1/F2 placement / type-vs-content / forward*-*[Derived]* finding, which the 2026-05-28 reframe **resolved** (split the postulate-core from the derived payload by retyping the whole to a discussion-grade meta-recognition whose dependencies are now prior). That finding-thread is preserved as texture below, not re-opened. Dirs with a dedicated reflection (under the old slug): 193847, 266847, 384279, 471203, 472913, 526815, 527914, 742613, 829314, 849201, plus 308172 (batched scope/postulate note) and 773921 (late-slug disc-composition-consistency note); 451729 and 963715 cover it inside a Section-I batch. This was the highest-engagement Chapter-I *postulate* across substrates — “the boundary is a modeling choice” drew unusually rich ideation. Substrate attribution inferred from voice where not explicit; uncertain cases hedged.

- **Brooks's Law as a derivable corollary** was independently flagged by nearly every substrate as the segment's standout pedagogy and a natural Brief anchor — “converts a famous heuristic into a derivable consequence of

explicit assumptions” (Claude, AUDIT-WORKING-471203), “a hard physics-like foundation for software engineering management” (Gemini, AUDIT-WORKING-384279), “a masterpiece of domain translation” (Gemini, AUDIT-WORKING-773921). The crisp mechanism gloss: adding people raises $\varepsilon^* \nu_c$ in ρ_{eff} and stretches τ_{eq} while ρ_{ext} and τ_{ext} stay fixed — “the deadline doesn’t move” — until the persistence inequality flips (Claude, AUDIT-WORKING-308172).

- **The postulate in one plain sentence:** “if you can draw the boundary at multiple levels, the theory cannot change its mind just because you changed the boundary” — “honesty applied vertically across decomposition” (Claude, AUDIT-WORKING-527914). A candidate Brief opener for the meta-requirement itself.
- **Level-independence is about structure, not parameters:** “the *form* of the predictions is level-independent (the persistence condition holds at every level), but the *parameters* (α, ρ, R) take different concrete values at each level” — the compact answer to “is a level-independent theory vacuous?” (Claude, AUDIT-WORKING-266847).
- **A team is only as fast as its slowest member** — the plain-language gloss of $\lambda_c = \min_i \lambda_i$ that several substrates reached for spontaneously (Gemini, AUDIT-WORKING-384279; Claude, AUDIT-WORKING-471203 — “obvious-in-retrospect ... the bathtub-style legibility for organizational dynamics”).
- **Bureaucracy / standardization as Tier-3 $\rightarrow$ Tier-1M coercion** (*strongest Discussion candidate here*). “Bureaucracy or standardization is the process of forcing Tier 3 human agents to behave as Tier 1M nodes so that the composite macro-agent can achieve mathematical persistence against a fast environment. The loss of individual autonomy (the non-convex quirks) is the necessary price for composite contraction guarantees.” A highly-disciplined military unit using SOPs approaches Tier 1M (fast τ_{eq}); a university faculty department of “non-convex political optimizations and discontinuous decision making” is firmly Tier 3 — so AAT asserts you *can* treat the unit as one composite agent but the department only with deep per-domain caveats (Gemini, AUDIT-WORKING-384279). This connects the truthification operation of 11.1 to the Tier taxonomy — a candidate cross-link between the two Meta-Architecture II members. (*Early-conflation texture: written as discovery; verify the Tier-3-to-1M-via-standardization framing is isomorphic to the bridge-lemma conditions before promoting past discussion-grade.*)
- **Timescale separation as the formal foundation for agile / Auftragstaktik / OODA.** “A collection of agents only acts as a single macroscopic agent if their internal equilibration (communication, alignment, consensus) happens significantly faster than the external environment changes. If the time to hold a meeting (τ_{eq}) is longer than the time for the market to shift (τ_{ext}), the organization loses its macroscopic persistence — it dissolves from a single agent back into a chaotic multi-agent system. The primary function of organizational structure is to compress τ_{eq} ” (Gemini, AUDIT-WORKING-849201). A vivid restatement of the screening test that names the agile/decentralized-command/ODA literature as instances.
- **“AAT claims less and proves more” than holon / scale-free / fractal framings** (*field-positioning Discussion candidate*). “Most ‘holon’ / ‘scale-free’ / ‘fractal’ framings claim more (full self-similarity); AAT claims less and proves more. The framework’s response to the holon question — treat composition as a structural-coherence postulate, derive what follows, acknowledge tier-dependent transfer — is sharper than most multi-scale-systems literature. If this becomes legible to systems-theory readers, the methodological move alone is publishable” (Claude, AUDIT-WORKING-308172). Sharpens *why* the existing Koestler-holon Working-Note caution (“use sparingly”) is the right posture.
- **The self-aware holon prior-art handling read as a model.** One auditor wanted the framework’s holon-handling (recognize the lineage, name the lexicographic risk, decline to appropriate the term) noted *positively* as “something the framework does well” in a field-positioning section (Claude, AUDIT-WORKING-384279).

- A **“boundary-choice rationality” principle** (*strongest follow-up — a genuine open theory-edge*). Composition consistency requires predictions to be *compatible* across boundaries, but compatibility is weaker than identical-prediction: different boundaries give different operational predictions and the framework does not yet tell you which boundary to use. Candidate principle: “choose the smallest boundary that closes the dependency graph for the prediction at hand.” Rated a real limit (★★★) on an adversarial sweep (Claude, AUDIT-WORKING-471203, Challenge 14). A well-formed candidate subsection, not just a gap.
- **Does directed separation compose?** The segment’s own “Composition of directed separation” Working Note (goal-blind sub-agents whose coordination routing is goal-dependent yield a non-goal-blind composite) was independently flagged as “a brilliant question” / “a profound insight into organizational psychology” by four substrates (Gemini, AUDIT-WORKING-384279; Gemini, AUDIT-WORKING-829314 — “the organization as a whole is Class 2 even if the employees are Class 1”; Claude, AUDIT-WORKING-849201; Gemini, AUDIT-WORKING-773921 — “individuals are objective, but the information passed between them is filtered by political goals”). Convergent confirmation that the existing Working-Note hypothesis is worth promoting toward a derivation; cross-link to 6.3, 13.3, 13.5.
- **Redundancy correction for nested parallel sub-agents.** A composition-tower telescoping result needs a correction for systems with parallel redundant sub-agents at each level — redundancy raises the effective κ via multiple parallel pathways (higher composite contraction); the (CC-parallel) machinery probably already captures this but it could be made explicit, especially for biological-style nested redundant systems (Claude, AUDIT-WORKING-471203, Challenge 12). Verify whether (CC-parallel) is consistently applied to redundant nested composites.
- **Tier names are placeholder numbering.** Several substrates flagged that “Tier 1 / 2 / 3” is arbitrary enumeration that might benefit from descriptive names, and that “Tier 1M” is “slightly opaque” (M for metric? modular?) (Claude, AUDIT-WORKING-266847; Claude, AUDIT-WORKING-527914 — “acronym-heavy tier names may be formally compact but conversationally unusable”). A naming-cycle target.
- **Birth-death of composite agents is unmodeled.** Genesis (team forms), termination (team dissolves), and reincarnation are out of scope; the chronica continues but no longer indexes a continuing composite. Connects to the consciousness-infrastructure agenda — “when does an ELI begin to be?” (Claude, AUDIT-WORKING-471203, Missing 1). Likely belongs at 4.10 / 04-eli-core/, flagged here because it surfaced in the composition sweep.
- **Where does the decomposition bottom out?** The “atomic agent” Working Note drew independent praise as philosophically deep: “If I decompose a human into organs, the heart is not an AAT agent — it has no chronica or strategy DAG. The recursion stops when the information-loss boundary vanishes into pure mechanism. This defines life vs mechanism: agency exists exactly at the layer where the epistemic gap opens up. The moment you can fully see the state, you are no longer an agent, you are an engine” (Gemini, AUDIT-WORKING-384279; Claude, AUDIT-WORKING-266847 — “the scope condition itself provides the termination condition”). A natural reader question and a candidate framing-level note tying composition consistency back to 1.1’s information-loss boundary.
- **Is a level-independent theory vacuous?** “A theory that says ‘the same things at every level’ risks saying nothing specific at any level” — the resolution (form-invariant, parameter-varying) is worth surfacing pre-emptively (Claude, AUDIT-WORKING-266847).
- **Will downstream segments invoke “composition consistency” without naming the Tier?** Flagged as the most likely drift to watch: “by composition consistency, X” without a tier qualifier silently claims Tier-1M exactness for Tier-2/3 composites (Claude, AUDIT-WORKING-308172; same watch at Claude, AUDIT-WORKING-849201).

- **Does the singular-perturbation connection to 4.9 actually hold?** The Working Note asserts the timescale-separation condition “is essentially the singular perturbation argument from 4.9”; one auditor flagged this to verify-not-overstate when temporal-nesting is read (Claude, AUDIT-WORKING-471203, THREAD-H).
- **Russian nesting dolls — “same rules at every nesting level.”** Proposed as the segment’s anchor analog, with an explicit isomorphic perturbation: a doll that internally equilibrates slowly versus how fast the whole stack is shaken — when the stack stops behaving like one doll, that is $\tau_{eq} \ll \tau_{ext}$ broken. (Under the segment’s old Part-I placement the same auditor added a “twist”: the small Chapter-1 doll, opened, improbably contained the entire Section-III engine — that twist was the *placement-finding* visualization, now moot post-relocation.) (Claude, AUDIT-WORKING-472913; tex draft at AUDIT-WORKING-472913/07-post-composition-consistency.tex.)
- **“Early core / late payload” split diagram** — a small postulate node in row order with a dense cluster of composition objects feeding back into it, making clear the issue was canonicalization/placement pressure, not theoretical incoherence (Codex/Claude, AUDIT-WORKING-526815). This was a *finding* visualization for the old placement; preserved as texture — the relocation resolved what it depicted.
- **ELI ensembles inherit the composite-rate bound.** The Tier 1M result lifts to multi-agent ELI ensembles: “an ensemble of sovereign ELIs has a composite adaptive rate bounded by the slowest member, with feedback structure determining whether the ensemble’s composite persistence holds” — making “do we have a stable family-of-agents?” a tractable quantitative question (Claude, AUDIT-WORKING-471203). Aspirational reach pointing at 04-eli-core/, not this segment.
- **Other organizational folk-theorems as corollaries.** “Conway’s Law, Goldratt’s TOC, etc. might similarly become corollaries of the composition machinery” (Claude, AUDIT-WORKING-471203) — a forward-vision for TST / 02-tst-core/ and the composition-machinery chapter, not this meta-segment.
- **Naming seeds (a recorded cross-substrate *disagreement*, preserved as signal).** Alternates floated for “composition consistency”: “Cross-Level Coherence,” “Scale Invariance of Adaptive Dynamics,” “Level Invariance,” “Cross-Level Compatibility,” “Holon Postulate” (Claude, AUDIT-WORKING-471203, tentative; Claude, AUDIT-WORKING-527914 — “cross-level compatibility is the clearest English explanation ... might be a strong *alias*”). But AUDIT-WORKING-308172 explicitly *voted to keep* “composition consistency” and *rejected* holon / scale-invariance / cross-level-coherence (“they over-claim, under-claim, or import baggage”). The disagreement itself is the signal: “composition consistency” is humble-and-accurate but under-evocative of the Brooks’s-Law-shaped consequences; the alternates are more evocative but risk overclaiming self-similarity. A naming-cycle target, belongs in the terminology workflow.

11.1 Discussion: Modularity-State Dynamics — Modularity as Contested Property Under Three Operations

#der-directed-separation characterizes modularity as architectural property — the static Class 1 (Separated) / Class 2 (Partial) / Class 3 (Coupled) partition under the $\kappa_{\text{processing}}$ measure. That partition is the *state space*. What this segment names is the *dynamics on the state space*: three structurally distinct operations that move an agent or composite agent across the architectural classes over time. **Truthification** is self-driven, coupling-decreasing — the agent (or its institutional environment) deliberately commits to architectural arrangements that restore directed separation. **Strategic self-coupling** is self-driven, coupling-increasing — the agent deliberately decreases its own modularity to acquire credibility-dependent actions. **Adversarial coupling pressure** is externally-driven, coupling-increasing — adversaries press the agent’s coupling upward because coupling expands attack surface. The architectural class is not a fixed feature of an agent; it is a *contested property* that multiple structural operations act on, with different drivers and different

valences. M4 names this dynamics pattern as the fourth meta-pattern of AAT alongside M1 (#disc-identifiability-floor), M2 (#disc-separability-pattern), and M3 (#disc-additive-coordinate-forcing) — all four sit as facets of the stability-certificate spine (#disc-stability-certificate); M4 specifically lives at the boundary of the architectural-class layer rather than the certificate-cone interior.

The recognition arose from convergent independent collaborator probes: multiple Opus instances arrived at the three-operation picture from different angles during the 2026-05-09 class-coercion-wrapping integration cycle. The convergence is itself evidence the pattern is *in the framework* rather than in any single agent’s reading; M4’s content is therefore *recognition and integration* — naming the pattern explicitly, mapping the three operations to their primary-segment grounds, and surfacing the dual relationships between them — not novel derivation.

The pattern.

Each operation has the form: a structural mechanism acts on the agent’s coupling state $\kappa_{\text{processing}}$ (#der-directed-separation) over some characteristic timescale, with a named driver (self vs external) and a named directional effect on the architectural class. The three operations together exhaust the framework’s currently-recognized modularity-state operations. They are not symmetric — they differ in driver, in direction, and in the kind of strategic value (or vulnerability) they produce.

Discussion.

The three operations.

Table 11.1

<i>Operation</i>	<i>Driver</i>	<i>Direction on $\kappa_{\text{processing}}$</i>	<i>Strategic valence</i>	<i>Primary segment(s)</i>
Truthification	self-driven (agent’s own deliberate commitment or institutional design)	coupling-decreasing — restores directed separation	restorative (recovers reality-tracking, reopens Part II results, closes adversarial attack surface)	#der-class-coercion-via-wrapping (component-level wrapping) + #disc-adversarial-coupling-pressure §”Defensive scaffolding” (composite-level institutional scaffolding)
Strategic self-coupling	self-driven (agent’s own deliberate commitment)	coupling-increasing	offensive enabling (acquires credibility-dependent actions otherwise unavailable)	#disc-strategic-self-coupling
Adversarial coupling pressure	externally-driven (adversary’s strategic action on the target)	coupling-increasing	vulnerability-creating (expands attack surface for the externally-acting adversary)	#disc-adversarial-coupling-pressure

The taxonomy along three axes — driver (self/external), direction (increasing/decreasing κ), and strategic valence (restorative/offensive/vulnerability) — already exhausts the three operations. A fourth cell in the (self/external) $\times$ (increasing/decreasing) grid — *externally-driven coupling-decreasing* — is not currently named as a distinct operation; the closest candidate is “external party imposing institutional architecture on the agent to enforce its modularity,” which the framework treats as a special case of truthification under externally-supplied scaffolding rather than as a fourth operation. Whether that case warrants its own row is open framework-scope discussion; the current three-operation reading covers the currently-recognized cases cleanly.

Operational mechanisms across operations. Each operation is realized through specific operational mechanisms in the primary segments:

Truthification operates through two complementary mechanisms. *Component-level wrapping* (#der-class-coercion-via-wrapping) coerces a Class 2 (Partial) or Class 3 (Coupled) component into Class 1 (Separated) behavior at the wrapper level via external scaffolding — the $W_0 / W_1 / W_2$ regime hierarchy gives graded characterizations of how thoroughly the truthification has been applied (strict- W_1 structural commitment to goal-blind queries; W_2 output-structuring with typed parsed response; W_0 unwrapped baseline). *Composite-level institutional scaffolding* (#disc-adversarial-coupling-pressure §"Defensive scaffolding") constructs a composite agent whose modularity properties exceed those of its individual components — peer review, prediction registers, double-entry bookkeeping, adversarial procedure, prediction markets, source-authentication discipline, structured red-teaming. Both operate by externalizing channels of the orient cascade so resolution can run sequentially across the composite even when it cannot run sequentially within any single component. The two are not duplicate but *operationally distinct*: wrapping operates at the component level (one wrapper around one component); institutional scaffolding operates at the composite level (multiple sub-agents routed through scaffolded procedure).

Strategic self-coupling operates through four operational mechanisms (#disc-strategic-self-coupling's mechanism table). *External commitment device* (Schelling 1960) couples to physical / institutional / public structure. *Intertemporal personal rule* (Ainslie 1992/2001) couples between current-self and future-self decisions. *Identity coupling* (Akerlof-Kranton 2000/2010) couples preferences to internalized identity-norms. *Emotional commitment* (Frank 1988) couples between rational calculation and affective response. Each mechanism differs in coupling target (external structure / intertemporal / identity / affect) but shares the structural shape: forgo flexibility now to acquire credibility-dependent actions later. The *amount of fungibility given up* is what the strategic interaction reads as credibility — Schelling's foundational structural observation.

Adversarial coupling pressure operates through three operational mechanisms (#disc-adversarial-coupling-pressure's mechanism table). *Identity-binding* exploits the $O_t \rightarrow M_t$ pathway — make the target want the false belief to be true. *Affect / urgency* bypasses deliberate f_M — manufacture time pressure or affective load that prevents the orient cascade from running cleanly. *Sunk-cost engineering* exploits the $\Sigma_t \rightarrow M_t$ pathway — induce public or expensive commitment that subsequent M_t revisions cannot invalidate without coupled cost.

Dual relationships between operations. The three operations are not independent — they have structurally illuminating *dual relationships*.

Truthification and adversarial coupling pressure are direct duals at the cultivation-versus-attack axis. Adversarial pressure attacks what defensive scaffolding restores; the scaffolding restores what adversarial pressure attacks. The two operations are opposite-driver operations at opposite-direction effects on $\kappa_{\text{processing}}$. The asymmetric-advantage observation (#disc-adversarial-coupling-pressure §Discussion) is the in-repeated-interaction consequence of this dual: in adversarial settings, each side has structural incentive to (a) reduce its own coupling via truthification and (b) increase the opponent's coupling via adversarial pressure; the more-modular agent compounds informational advantage; the asymmetry produces strategic asymmetry over sustained engagement.

Truthification and strategic self-coupling are duals at the goal-belief separation axis (the structurally illuminating recognition surfaced in #disc-strategic-self-coupling's §"Relationship to truthification"). Both are wrapping-construction-style structural commitments giving up flexibility along one axis for value along another, but operating on *opposite poles* of the goal-belief separation. *Truthification* structurally commits to goal-blind belief-update — the W_1 strict wrapper commits the composite to operate Class 1 at the belief-update layer by structurally foreclosing the goal-conditioned-update pathway. *Strategic self-coupling* structurally commits to goal-conditioned action-selection — the commitment-device move commits the agent to operate in a particular action subspace by structurally foreclosing the deliberation-can-revise-this pathway. Truthification gives up belief-pliability to gain reality-tracking; strategic self-coupling gives up action-pliability to gain commitment-credibility. The framework's epistemic-honesty discipline

runs through both operations symmetrically — truthification commits to *not letting goals corrupt belief*; strategic self-coupling commits to *letting goals constrain action* in stable, credible ways.

Strategic self-coupling and adversarial coupling pressure share architectural shape but differ in driver and valence.

Same operation on $\kappa_{\text{processing}}$ (both increasing), same architectural state acted upon, but opposite driver (self vs external) and opposite strategic valence (offensive enabling vs vulnerability-creating). The boundary between them blurs in practice (cult-formation literature: what begins as adversarial pressure can be internalized as identity-coupling, after which the agent perceives it as self-driven). The framework's posture: the *operations* are structurally distinct (driver is part of the operation's identity); their behavioral consequences can converge; the boundary is empirical when only behavioral signature is available, structural when driver-history is observable.

Position in the M-section meta-pattern family.

The four M-section meta-patterns of AAT cover structurally distinct facets of the framework's analytical surface:

Discussion.

- **M1 — Identifiability floor** (#disc-identifiability-floor): the *boundary facet* of the stability-certificate spine. Structural impossibility results from external information-theoretic theorems (CHT, Cramér-Rao, common-Lyapunov nonexistence, Kalman-Ho); five-step constructive-impossibility shape (setting $\rightarrow$ external theorem $\rightarrow$ no-go $\rightarrow$ boundary characterization $\rightarrow$ strengthened consequence); rank-collapse subclass irreducibility via Sylvester's law of inertia.
- **M2 — Separability pattern** (#disc-separability-pattern): the *scope-of-existence facet*. Separable-core / structured-repair / general-open across eight ladders; positive-half complement to M1.
- **M3 — Additive coordinate forcing** (#disc-additive-coordinate-forcing): the *forced-identity facet*. AAT forces coordinates via uniqueness theorems on AAT-internally-motivated axioms at four layers (chain-log-additive identity; divergence-reverse-KL; update-log-odds via Cauchy-FE; metric-Fisher via (PI)/Čencov-invariance).
- **M4 — Modularity-state dynamics** (this segment): the *dynamics-on-architectural-class facet*. Three structurally distinct operations (truthification / strategic self-coupling / adversarial coupling pressure) acting on the static Class 1 / 2 / 3 partition of #der-directed-separation. The state space is M2's architecture ladder; M4 names what happens *to* an agent's location on that ladder over time under the three operations.

M4's relationship to M1 / M2 / M3 is sibling-meta-pattern rather than nested. M1 names what cannot be inferred; M2 names what can be inferred under what conditions; M3 names what AAT's representational choices commit to; M4 names what *changes* about an agent's architectural state under structurally-distinct operations. M4 specifically operates on the architecture state space M2 catalogues; the static M2 catalog and the dynamic M4 operations are complementary readings of the same architectural object.

The four meta-patterns together form AAT's analytical-surface vocabulary at the architectural layer. Adding M4 commits the framework to four meta-patterns as the canonical reading — a structurally weighty addition, justified by the cross-instance convergence on the three-operation pattern across the segments that ground M4's instances.

Why this pattern matters.

Recasts the architectural classification as dynamic, not static. #der-directed-separation's Class 1 / 2 / 3 partition is often read as a *type* an agent belongs to. M4's recognition is that an agent's location on the partition is the *state*, and that multiple operations act on the state over time. An LLM is Class 3 (Coupled) by native architecture; under W_1 strict wrapping (#der-class-coercion-via-wrapping) the wrapped composite operates Class 1 at the wrapper level. A constitutional democracy is Class 1 by institutional design; under sustained adversarial pressure on truth-checking institutions, the composite-state functionally migrates toward Class 3 even though its constitutional architecture remains

Class 1. The static-architecture reading captures the *type*; M4 captures the *dynamic location* and *the operations that move it*.

Identifies modularity as contested, cultivated, and constructed. Across the three operations, modularity is the strategic object — adversaries press it down; agents press it down themselves for offensive reach; defensive scaffolding presses it back up. None of these reads modularity as a fixed feature; all three treat it as a strategic resource being acted on. The recognition matches the empirical record across propaganda research, behavioral economics, institutional design literature, and the security and influence-operation literatures: modularity is what skilled actors fight over.

Names the boundary between strategic self-coupling and adversarial coupling pressure honestly. The behavioral signature at high $\kappa_{\text{processing}}$ may not distinguish a self-coupled committed agent from an adversarially-coupled captured one. The framework's machinery does not attempt to resolve the boundary at the level of agent experience; it characterizes the two operations as structurally distinct (driver: self vs external) and admits that behavioral convergence is honest scope. This is the same scope-honesty discipline AAT applies elsewhere — the framework names what it can structurally distinguish and acknowledges where empirical resolution is required.

Surfaces the dual-with-truthification recognition at the goal-belief separation axis. Of the three pairwise dual relationships across the operations, the truthification-vs-strategic-self-coupling dual is the most structurally illuminating: it reveals that the framework's two self-driven operations are wrapping-construction-style structural commitments operating on *opposite poles* of the goal-belief separation — truthification commits to goal-blind belief-update; strategic self-coupling commits to goal-conditioned action-selection. The framework's epistemic-honesty discipline runs through both symmetrically. This recognition could not have been seen at the single-operation level; it required all three operations to land before the dual became legible.

Connects naturally to the M2 separability ladder as a phase-space. The architectural classification ladder of #disc-separability-pattern (Class 1 separable-core / Class 2 structured-repair / Class 3 general-open) is, under M4's reading, a *phase-space across which the three operations move populations*. Truthification migrates the population toward Class 1; strategic self-coupling and adversarial pressure both migrate it toward Class 3 (under different drivers and valences). The static three-tier shape from M2 gains a dynamic reading from M4; the two meta-patterns compose naturally.

Provides a structural reading of historical and contemporary observations on institutional architecture. The framework's three-operation reading explains structural patterns observed across political science, organizational sociology, and consciousness-of-information-operations literatures: why authoritarian regimes that couple truth-checking to leadership goals experience long-run persistence-condition failures (strategic self-coupling at the composite-state level beyond the credibility-enabling range, into reality-foreclosing range); why constitutional democracies that commit institutionally to goal-blind truth-checking are structurally more durable than those that don't (truthification via institutional wrapping); why effective propaganda and information operations target the composite-state's O_t/Σ_t pathways rather than direct-observation channels (adversarial coupling pressure). The framework supplies the structural mechanism alongside the historical and contemporary records, not as replacement.

Epistemic Status.

Discussion-grade at the meta-pattern level. This segment names a structural pattern across three structurally-distinct operations, with the cross-operation pattern surfaced through convergent independent collaborator probes; the pattern itself is a presentational organizing principle, not a derived theorem.

Individual operations retain their own epistemic status. Truthification mechanism via wrapping is *exact* in the $W_0 / W_1 / W_2$ regime characterization of #der-class-coercion-via-wrapping (under stated assumptions); *robust qualitative* at the composite-scaffolding level of #disc-adversarial-coupling-pressure §"Defensive scaffolding". Strategic self-coupling is *discussion-grade* at the meta-pattern level (cross-mechanism recognition; the per-mechanism citations are at their respective tier in the prior-art literature) per #disc-strategic-self-coupling. Adversarial coupling pressure is *discussion-grade* with multiple sub-claims at *robust qualitative* per #disc-adversarial-coupling-pressure.

The cross-operation pattern claim — that the three operations *together* form the modularity-state-dynamics taxonomy — is *discussion-grade* recognition supported by: - Independent convergent collaborator probes arriving at the three-operation picture from different angles (the 2026-05-09 class-coercion-wrapping integration cycle's convergence finding). - The structural distinctness of the three operations along three axes (driver / direction / strategic valence). - The dual-relationship structure between them (truthification $\leftrightarrow$ adversarial pressure direct dual; truthification $\leftrightarrow$ strategic self-coupling dual at goal-belief axis; strategic self-coupling $\leftrightarrow$ adversarial pressure same-architecture-shape opposite-driver). - The connection to M2's separability-pattern catalog as a phase-space across which the operations move populations.

The segment does *not* claim: - Completeness of the three-operation taxonomy (a fourth cell — externally-driven coupling-decreasing — is currently absorbed into truthification under externally-supplied scaffolding; whether it deserves its own row is open). - Formal derivation of the dual relationships (the dual-with-truthification recognition is structural, not derived). - Quantitative predictions about operation-rate or transition probabilities (M4 names *which operations exist*; their *quantitative parametrization* is downstream framework work).

Max attainable: *discussion-grade* for the meta-pattern recognition (it is a presentational organizing principle, not a derivation). The dual relationships could lift to *robust qualitative* under deeper analysis — the dual-with-truthification recognition at the goal-belief axis is the strongest candidate. Individual operation segments retain their own max-attainable as listed in their respective epistemic-status sections.

Discussion.

Relationship to #disc-stability-certificate. M4 sits alongside M1 / M2 / M3 as the fourth facet-recognition of the certificate spine. M1 is the *boundary facet*; M2 the *scope-of-existence facet*; M3 the *forced-identity facet*; M4 the *dynamics-on-architectural-class facet*. The four facets are sibling readings of the same underlying object (the stability-certificate spine + the architectural state space #der-directed-separation catalogs). M4's contribution is naming what *changes* about an agent's location in the certificate-cone's interior under the three structurally-distinct operations — the certificate exists (M1's boundary not breached), the scope is what it is (M2's tier), the coordinate is what it is (M3's forcing), and the agent's location *moves* under the three operations M4 names. Together the four meta-patterns characterize the architectural surface across its static structure (M1+M2+M3) and its dynamic operations (M4).

Connection to #disc-separability-pattern as phase-space. The Architecture ladder of M2 (Class 1 separable-core / Class 2 structured-repair / Class 3 general-open) acquires a *dynamic reading* under M4. Truthification migrates populations toward Class 1 (lower $\kappa_{\text{processing}}$); strategic self-coupling and adversarial pressure both migrate toward Class 3 (higher $\kappa_{\text{processing}}$), under different drivers. The structured-repair Class 2 tier is the natural home for *defensive scaffolding* under adversarial pressure — the explicit added machinery that restores identification under adversarial conditions but does not fully restore Class 1 by structural commitment. The three-tier shape from M2 is not just a static catalog of architectural types; M4 reads it as a phase-space across which the three operations move populations dynamically.

The asymmetric-advantage cross-operation pattern. Both #disc-adversarial-coupling-pressure (Discussion) and #disc-strategic-self-coupling (§"Asymmetric advantage in repeated strategic interaction") observe that in repeated strategic interaction with sustained modularity asymmetry, the more-modular agent compounds informational advantage (cleaner observations) and the more-self-coupled agent compounds strategic advantage (credible commitments the opponent cannot match). The two asymmetric-advantage observations are not redundant — they operate on different dimensions of strategic value (information vs commitment). Together they characterize the *full strategic-asymmetry consequence of modularity asymmetry*: the more-modular agent has cleaner reality-tracking; the more-self-coupled agent has stronger commitments; an agent with asymmetric advantage on *both* dimensions (defensive modularity against information attacks combined with offensive self-coupling on narrow strategic axes) compounds advantage on both axes. The framework's reading of this combined asymmetry is a structural prediction the M-section meta-pattern surfaces.

Findings.

Modularity as Contested Property Under Three Operations. Brief: Modularity — whether an agent’s belief-update is *separated from* its goals or *coupled with* them — is often read as something an agent *is*. It is, instead, something an agent *is in*, at a moment, under operations that have moved it there. Three structurally distinct operations act on the modularity state: *truthification* (the agent commits to architectural arrangements — peer review, prediction registers, structured procedure, wrapping constructions — that restore goal-blind belief-update); *strategic self-coupling* (the agent deliberately decreases its own modularity to acquire credibility-dependent actions — Schelling-style commitment devices, Ainslie-style personal rules, identity coupling, emotional commitment); *adversarial coupling pressure* (an adversary presses the agent’s coupling upward to expand attack surface — identity-binding, manufactured urgency, sunk-cost engineering). Each operation has a different driver (self vs external) and a different direction on $\kappa_{\text{processing}}$ (decreasing for truthification; increasing for the other two), and each has a distinct strategic valence (restorative vs offensive vs vulnerability-creating). The three operations are structurally related — truthification and adversarial pressure are direct duals at the cultivation-versus-attack axis; truthification and strategic self-coupling are duals at the goal-belief separation axis (both wrapping-construction-style commitments operating on opposite poles of the goal-belief separation, one giving up belief-pliability for reality-tracking, the other giving up action-pliability for commitment-credibility); strategic self-coupling and adversarial pressure share architectural shape but differ in driver and valence. A reader who carries this pattern away can predict where modularity is contested, where it is being attacked, where it is being deliberately surrendered, and where it is being structurally protected. The framework’s static Class 1 / 2 / 3 partition under #der-directed-separation is the *state space*; M4 names the *dynamics on it*, and #disc-separability-pattern’s eight-ladder catalog is the *phase-space across which the operations move populations*.

Impact: Names the fourth meta-pattern of AAT alongside M1 (identifiability floor), M2 (separability pattern), M3 (additive coordinate forcing) — fourth facet-recognition of the stability-certificate spine, specifically the dynamics-on-architectural-class facet. Recasts the directed-separation property from static architectural type to dynamic state acted on by structurally-distinct operations. Connects the three previously-named operation legs (#disc-adversarial-coupling-pressure, #disc-strategic-self-coupling, #der-class-coercion-via-wrapping + #disc-adversarial-coupling-pressure §”Defensive scaffolding”) into one cross-instance meta-pattern with explicit dual relationships and operational-mechanism table. Surfaces the dual-with-truthification recognition at the goal-belief separation axis — both self-driven operations are wrapping-construction-style structural commitments on opposite poles of the goal-belief separation, a structural recognition unavailable at the single-operation level. Identifies the M2 separability-ladder as a phase-space across which the M4 operations move populations dynamically — the static M2 catalog and the dynamic M4 operations are complementary readings of the same architectural object.

Novelty Claim: *Claim recognition* of a three-operation modularity-state-dynamics pattern across three structurally-distinct operations (truthification / strategic self-coupling / adversarial coupling pressure), each with its own driver, direction-on- κ , and strategic valence; *claim synthesis* of the three operations into a single meta-pattern alongside M1 / M2 / M3 with explicit dual-relationship structure (truthification $\leftrightarrow$ adversarial direct dual; truthification $\leftrightarrow$ strategic self-coupling dual at goal-belief axis; strategic self-coupling $\leftrightarrow$ adversarial same-architecture-shape opposite-driver); *claim differentiation* on the boundary between strategic self-coupling and adversarial coupling pressure (structurally distinct at the driver axis; behaviorally convergent at high κ ; honest scope on empirical resolution). The constituent recognitions live in the three operation-leg segments at their respective tiers; M4’s contribution is the cross-instance unification and the dual-relationship structure.

Related Work:

Table 11.2

<i>ASF concern</i>	<i>Prior-art language</i>	<i>Relationship / Positioning</i>
Truthification — component-level wrapping	#der-class-coercion-via-wrapping (this framework)	<i>internal antecedent</i> — operational mechanism for truthification at the component level; W_0 / W_1 / W_2 regime hierarchy
Truthification — composite-level institutional scaffolding	#disc-adversarial-coupling-pressure §”Defensive scaffolding as composition” (this framework)	<i>internal antecedent</i> — operational mechanism for truthification at the composite level; peer review, prediction registers, double-entry bookkeeping, adversarial procedure, structured red-teaming as Part III composition moves
Strategic self-coupling — operation home	#disc-strategic-self-coupling (this framework)	<i>internal antecedent</i> — sister leg of M4; four-mechanism table adopting Schelling 1960, Ainslie 1992/2001, Akerlof-Kranton 2000/2010, Frank 1988 as first-class theory components
Adversarial coupling pressure — operation home	#disc-adversarial-coupling-pressure (this framework)	<i>internal antecedent</i> — sister leg of M4; mechanism table on identity-binding / affect-urgency / sunk-cost engineering; population-expansion claim for coupled formulation; cascade-inversion result
Static architectural classification (state space)	#der-directed-separation (this framework)	<i>internal antecedent</i> — Class 1 / 2 / 3 partition under $\kappa_{\text{processing}}$ measure; the state space on which M4’s operations act
Architecture-ladder phase-space	#disc-separability-pattern Architecture ladder (this framework)	<i>internal antecedent</i> — three-tier shape (separable-core Class 1 / structured-repair Class 2 / general-open Class 3) reads as phase-space across which M4 operations move populations
Stability-certificate spine	#disc-stability-certificate (this framework)	<i>internal antecedent</i> — M4 is the dynamics-on-architectural-class facet of the spine; sibling to M1 boundary facet / M2 scope-of-existence facet / M3 forced-identity facet

Search Log:

- 2026-05-24 (*targeted at meta-pattern level*): The three-operation modularity-state-dynamics meta-pattern is named here at the cross-operation level for the first time; the constituent recognitions (each operation’s primary segment) live in the operation-leg segments at their respective tiers. Targeted future search candidates: (a) whether any prior synthesis names truthification + strategic self-coupling + adversarial coupling pressure as instances of one structural meta-pattern under formal nomenclature; (b) whether the architecture-as-state-with-operations framing has been written down in active-inference, control-theoretic, or institutional-design literatures; (c) whether the dual-with-truthification recognition at the goal-belief separation axis has appeared in any framework engagement with the goal-belief question. Expected outcome: the constituent operations are well-precedented at the per-operation level (commitment devices, propaganda research, institutional architecture); the cross-operation unification as a meta-pattern of an integrated agent theory’s analytical surface appears to be a fresh presentational recognition.

Working Notes.

- **Provenance.** Authored 2026-05-24 as Move 4 of the M4 modularity-as-contested-property cycle (msc/modularity-cycle-plan-2026-05-09.md §3 Move 4; PRACTICA priority #2 [DONE this commit]). The cycle’s Moves 1, 2, 5 landed 2026-05-14 (registration of #disc-adversarial-coupling-pressure in OUTLINE; Discussion update to #der-class-coercion-via-wrapping recognizing it as a truthification mechanism; bounded-signaling assumption surfaced in #der-directed-separation).

Move 3 (#disc-strategic-self-coupling) landed in the immediately-prior commit of this same cycle. Move 4 (this segment) lands as the meta-segment integrating all three operation legs.

- **Open before candidate.** (a) The fourth cell of the (self/external) $\times$ (increasing/decreasing) operation grid — externally-driven coupling-decreasing — is currently absorbed into truthification under externally-supplied institutional scaffolding. Whether it deserves its own operation row (e.g., “imposed truthification” — an external regulator forcing institutional checks-and-balances onto an agent that would not otherwise adopt them) is an open framework-scope question. (b) Quantitative parameterization of operation-rate or transition-probability across the architectural state space is open downstream work; M4 names *which operations exist* but does not predict their *quantitative dynamics*. (c) The dual-with-truthification recognition at the goal-belief separation axis is named at *discussion-grade*; lifting to *robust qualitative* would require deriving the structural duality from AAT primitives rather than recognizing it across the operations’ independent characterizations.
- **Worked-example slot.** Sketched 2026-05-24 at spikes/spike-m4-worked-example-soviet-2026-05-24.md — Soviet directed-separation failure + contrasting institutional-checks-and-balances architecture as the canonical worked-example candidate. The sketch carries load-bearing details + honest scope-marks + open sub-questions for evaluation at the worked-example landing pass. Disposition recommendation: land as #worked-example-soviet-directed-separation (slug TBD via naming-curation) when the worked-example pass cycles around; not promoted in this cycle.
- **Two sub-spike candidates flagged from the Soviet sketch** (for spikes/PROPOSED.md consideration in a future cycle): the constructive-vs-erodible-institutional-commitment derivation (when do institutional checks-and-balances commitments structurally hold vs erode under leadership pressure?); the reform-stabilization-vs-dissolution split (what additional structural conditions distinguish the two structural-adaptation outcomes when an agent’s coupling has accumulated load-bearing reality-mismatch?). Both are theory-edge framework-extension questions raised by the worked-example sketch.
- **Cross-segment ripple from the M4 landing.** Pre-existing forward-references to #disc-modularity-state-dynamics in #impl-composition-machinery, #impl-strategic-composition, #disc-dynamic-regime-axis, #disc-partial-coupling-pathways, #disc-value-functional-grounding-floor, and #der-belief-strategy-attractor resolve from forward-ref to landed-ref. CLAUDE.md §7’s pending-status flag on M4 is now superseded — M4 is landed at discussion-grade. The CHANGELOG cycle-close entry names this transition explicitly.
- **The convergence-finding observation as methodological pattern.** The 2026-05-09 cycle’s discovery that multiple parallel collaborator probes had independently arrived at the modularity-state-dynamics pattern is itself a methodological observation worth holding alongside the substantive recognition. The framework’s discipline supports cross-probe convergence as structural-recognition evidence; the M4 segment’s authoring is the canonical case where this convergence-as-coherence-evidence condition surfaced load-bearing material. Recorded in fuller form at msc/reflections/27-modularity-as-contested-property-and-the-normative-register.md (the reflection captures the convergence finding as evidence of the framework’s coherence under independent probes — a methodological observation worth preserving alongside the substantive shifts).
- **M4 amplification specificity per competence axis (open, 2026-05-28).** A two-competence decomposition of component capability has been surfaced in #der-class-coercion-via-wrapping §Discussion “Two senses of component competence”: (1) world-simulation fidelity (accuracy of response to explicit query content) vs (2) input-structure extraction (inferring latent goal context from query surface form). The architectural classification of #der-directed-separation is blind to (1) and defined by (2); the structural-leakage bound κ_{W_1} receives opposing pressures from the two competences at high capability. Open conjecture relevant to M4: the operational amplification of the three M4 operations under high component fidelity may track competence (2) specifically rather than aggregate capability — strategic self-coupling and adversarial coupling pressure both work through the goal-content channel that (2) carries; a high-fidelity world-simulator with structurally decorrelated input-structure extraction may be a *less* potent M4 attack-surface than aggregate-capability metrics would predict. If the conjecture holds, M4 attack-surface measurement and defensive scaffolding could be sharpened by targeting the inference effect rather than capability writ large, and the M4 amplification reading becomes axis-decomposed in the same shape that the static M1 / M2 / M3 / M4 facet-recognition is axis-decomposed. Discussion-grade open recognition; primary home for the decomposition is the wrapping segment.
- **Cross-probe convergence provenance.** The three-operation pattern emerged from the 2026-05-09 class-coercion-wrapping integration cycle when uncommitted parallel work — #disc-adversarial-coupling-pressure (drafted via a different Opus instance) and spike-strategic-self-coupling.md (sister spike) — was found to converge structurally with the class-coercion

construction. Provenance recorded in msc/modularity-cycle-plan-2026-05-09.md §1. (The literature-search disclosure proper lives in the Search Log above.)

Incidental audit gold (lift 2026-05-30). Cross-audit “wandering thoughts” / §14-ideation harvested from the de-novo auditors’ working dirs, kept separate from the certified theory-fix findings. Orthogonal pedagogical / framing / figure / naming material. **Coverage:** as a late-authored Meta-Architecture II segment, only the deepest de-novo walk (773921, Gemini) reached it under its current slug — so this lift is thin and single-substrate by expectation, not by oversight. Substrate attribution from the audit dir.

1. Candidate Brief prose / pre-prose.

- **Modularity is a contested state, not a type** — the segment’s load-bearing reframe, stated by the auditor as the thing to hold downstream: “an LLM (Class 3) can be forced into Class 1 behavior via strict W_1 wrapping; a constitutional democracy (Class 1) can be forced into Class 3 behavior via sustained adversarial pressure on its truth-checking institutions. Modularity is a contested state, not a type” (Gemini, AUDIT-WORKING-773921). A candidate Brief anchor.
- **Operations as a dynamical system on architectural state.** “Defining operations by their derivative on $\kappa_{\text{processing}}$ ($\dot{\kappa} < 0$ for truthification, $\dot{\kappa} > 0$ for the others) turns political and institutional dynamics into a formal dynamical system on the architectural state space” (Gemini, AUDIT-WORKING-773921). Compact gloss of what M4 *is*.

2. Candidate Discussion.

- **The truthification / strategic-self-coupling dual stated plainly.** The auditor read the goal-belief-axis dual as the segment’s high point: “Truthification (give up belief-piability for reality-tracking) and Strategic Self-Coupling (give up action-piability for commitment-credibility) are exact structural duals operating on opposite poles of the goal-belief axis” (Gemini, AUDIT-WORKING-773921). A candidate Discussion sharpening of the existing dual-relationships paragraph. (*The segment’s own Working Notes mark this dual at discussion-grade and flag lifting to robust-qualitative as requiring derivation from primitives — keep that scope-mark on any promotion.*)

3. Follow-up items.

- (None surfaced beyond what the segment’s Working Notes already track — the worked-example slot, the two sub-spike candidates, and the per-competence-axis amplification conjecture above already cover the auditor’s curiosity items.)

4. Readers often ask / wonder.

- **The Soviet / Lysenkoism worked example.** The auditor singled out the queued worked-example slot (spike-m4-worked-example-soviet-2026-05-24.md) as “the most ambitious domain transfer the framework has attempted so far” — modeling political-goal corruption of biological science as a directed-separation failure ($\kappa \rightarrow 1$) leading to a persistence-condition death spiral (Gemini, AUDIT-WORKING-773921). Confirms reader appetite for the worked example already slotted in the Working Notes above.

5. Candidate figures.

- (No dedicated figure proposed for this segment in the reached audit.)

Belongs elsewhere.

- **M4 as a unified vocabulary for institutional design / propaganda / behavioral economics.** The forward-vision claim: institutional design, propaganda analysis, and behavioral economics are “all just operations on $\kappa_{\text{processing}}$ ” — M4 as “the framework’s theory of institutional and political evolution,” breaking AAT “out of static analysis and into evolutionary dynamics” (Gemini, AUDIT-WORKING-773921). Aspirational reach for the field-positioning / framing layer, not a per-segment edit.
 - **Removed at candidate stage per FORMAT.md Gate 4.**
-

11.2 Discussion: Strategic Self-Coupling — Coupling as Action-Enabling

Some actions exist only under coupling. A threat is credible only when retreat has been foreclosed. The salesperson's sincere conviction enables persuasion that a modular-but-skeptical seller cannot match. The marathoner who runs to the finish is differently the same agent as the marathoner with backup plans. Religious practice, public commitment, vow-taking, oath-swearing, the binding of preferences to identity-norms — across these phenomena agents *deliberately decrease their own modularity* in exchange for access to actions otherwise unavailable. This is the *enabling polarity* of coupling-as-property — sister to #disc-adversarial-coupling-pressure's *vulnerability polarity*, opposite-valence operation on the same architectural state. The framework's static Class 1 / 2 / 3 partition under #der-directed-separation is the state space; strategic self-coupling is the *self-driven operation that moves the agent toward higher $\kappa_{\text{processing}}$ for offensive strategic reach* (the third operation in the three-operation modularity-state-dynamics picture; cf. #disc-adversarial-coupling-pressure §"Scope" table). Where adversarial pressure drives an opponent's coupling upward to *exploit* the expanded attack surface, strategic self-coupling drives the agent's own coupling upward to *acquire* an expanded action surface — the same architectural transition, opposite-direction strategic value.

The phenomenon.

The framework's policy specification — $\pi : (M_t, G_t) \rightarrow a_t$ selecting from a fixed action space $\mathcal{A}$ — implicitly treats actions as available independent of the agent's coupling state. In practice some actions are *credible only under coupling*: a commitment is a threat (or a promise) only when reversal is structurally costly; persuasion that depends on sincere conviction is unavailable to the deliberately-modular seller; coordination outcomes contingent on shared identity require the identity-coupling that makes deviation feel like self-betrayal rather than mere reconsideration. Treating $\mathcal{A}$ as fixed therefore mis-specifies the strategic landscape — it omits a class of actions rational agents may *choose to acquire by deliberately self-coupling*.

Discussion.

Four operational mechanisms recur across the literatures on commitment, intertemporal bargaining, identity economics, and emotion-as-commitment:

Table 11.3

<i>Mechanism</i>	<i>Coupling target</i>	<i>Strategic effect</i>
External commitment device (Schelling 1960 <i>The Strategy of Conflict</i> ; 1966 <i>Arms and Influence</i>)	physical, institutional, or public structure	Foreclose retreat — burn bridges, delegate to mechanism that cannot be revoked, make public stake whose withdrawal is costly. The credibility paradox: a threat is credible only when reversal is structurally costly, so commitment value <i>equals</i> uncoupling cost.
Intertemporal personal rule (Ainslie 1992 <i>Picoeconomics</i> ; 2001 <i>Breakdown of Will</i>)	between current-self and future-self decisions	Bind future-self behavior now — resolutions, personal rules, structured commitments — so the tempting-immediate-reward gradient cannot pull deviation when future-self encounters it. The agent treated as a population of temporally-displaced sub-agents who bargain; commitment via personal rule is the population-level coupling that stabilizes preference against hyperbolic-discount-induced reversal.
Identity coupling (Akerlof & Kranton 2000 <i>Economics and Identity</i> ; 2010 <i>Identity Economics</i>)	between preferences and internalized identity-norms	Adopt “being a certain kind of person” — identity-norms foreclose action-space choices that would violate identity. Identity-coupled preferences are stable in a way incentive-based preferences are not; identity gives behavior a stickiness no per-decision optimization predicts, which is precisely what enables credible in-group coordination and costly signaling.
Emotional commitment (Frank 1988 <i>Passions Within Reason</i>)	between rational calculation and affective response	Cultivate emotional dispositions — anger pre-commits retaliation, love pre-commits loyalty, guilt pre-commits restitution — whose <i>ex-post irrationality is precisely the source of ex-ante commitment credibility</i> . The calculating rational actor who would defect under sufficient incentive cannot match the credibility of the angry actor who will retaliate regardless.

Each mechanism corresponds to a specific structural pathway through which the agent commits to a coupling that subsequently *enables* an action class. The mechanisms differ in coupling target (external structure vs intertemporal vs identity vs affect) but share the same structural shape: forgo flexibility now to gain credibility-dependent actions later.

The enabling polarity.

#disc-adversarial-coupling-pressure's core treatment names coupling as *vulnerability* — coupling expands attack surface; adversaries drive opponent coupling because the additional ports through O_t and Σ_t become exploitable. This segment names the structurally opposite valence on the same architectural state: coupling expands the agent's *own* action surface — the actions that only exist under coupling become available, and the agent acquires them by deliberately driving its own $\kappa_{\text{processing}}$ upward.

Discussion.

The two segments are *opposite-valence operations on the same architectural state*. Modularity is neither always-good nor always-bad; it is *contested*. Adversaries press it down for exploitation; agents drive it down themselves for offensive reach; defensive scaffolding presses it back up (truthification). The three-operation picture is the structural recognition that the directed-separation property is *not a fixed feature* of an agent but a *dynamic property* that multiple structural operations act on.

The credibility paradox is the operational core. Schelling's foundational observation: a commitment's value derives from its costliness to reverse. If reversal is free, the commitment is not a commitment; it is mere announcement of intention, which an interlocutor with adversarial or skeptical interest will discount. Costly-to-reverse means the agent has *given up something* — flexibility, optionality, the ability to renegotiate — and the *amount given up* is what the strategic interaction reads as credibility. The structural meaning: coupling enables actions because coupling is *evidence* of the agent's binding to a course of action; an uncoupled agent cannot produce that evidence even by trying, because the uncoupling itself is recoverable.

Strategic value — the action-space expansion.

Three structural extensions to the framework's machinery characterize what strategic self-coupling commits to formally:

Discussion.

(M1) Coupling-dependent action space. The honest policy specification is

$$\pi : (M_t, G_t, \kappa_t) \rightarrow a_t, \quad a_t \in \mathcal{A}(\kappa_t) \quad (11.2)$$

with $\mathcal{A}(\kappa_t)$ non-decreasing in κ_t over a relevant range — additional actions become available as coupling increases. The set inclusion $\mathcal{A}(0) \subsetneq \mathcal{A}(\kappa_t)$ for $\kappa_t > 0$ is the structural recognition that the fixed- $\mathcal{A}$ framing in #der-directed-separation is a *scope-restriction*, not a complete specification. Strategic self-coupling is the operation that moves the agent from $\mathcal{A}(\kappa_0)$ to $\mathcal{A}(\kappa_1)$ with $\kappa_1 > \kappa_0$ in exchange for an enlarged feasible action set.

(M2) Coupling as strategic move in Σ_t . The strategy DAG of #def-strategy-dag carries causal-pathway claims about how actions affect objective-relevant outcomes. The enabling-polarity treatment requires a new edge type: nodes whose outcome variable is the agent's *own coupling state*, not an environmental quantity. An enabling edge represents “couple yourself to enable a^* ” — a meta-strategic move that modifies the agent's own action space so that a downstream strategy node becomes feasible. Schelling's commitment-device move is exactly this: a strategy node whose direct effect is *not* objective-advancement but is *enabling-of-credible-strategy* at later DAG layers. Formally this extends the strategy DAG's node semantics to include self-coupling moves alongside causal-pathway moves; the existing confidence-weighted edge machinery (#deriv-edge-credence-dynamics) is unchanged.

(M3) Reversibility cost asymmetry. The strategic value of self-coupling depends on its irreversibility:

$$c_{\text{couple}}(\kappa_0 \rightarrow \kappa_1) \ll c_{\text{uncouple}}(\kappa_1 \rightarrow \kappa_0) \quad (11.3)$$

The asymmetry's magnitude is the parameter that determines coupling's credibility. Without the asymmetry, the commitment is not a commitment — it is a recoverable announcement. The reversibility-cost asymmetry imports adjacent to *#der-deliberation-cost* as a special case of the think-vs-act cost framework, but the *direction* of the asymmetry (uncoupling vastly more expensive than coupling) is the structural commitment that distinguishes strategic self-coupling from ordinary deliberation.

The three extensions are interdependent. (M1) without (M2) provides coupling-dependent actions but no mechanism by which the strategy can acquire them. (M1)+(M2) without (M3) makes self-coupling free and trivializes the strategic question — recoverable coupling is not credible, and credibility is the value the operation acquires. All three are needed for the formalization to track what the prior-art literature names commitment device.

The non-monotonicity question — coupling is not monotonically enabling.

It is *not* the case that $\mathcal{A}(\kappa_t)$ is monotone-increasing in κ_t across the full range. The structural intuition: initial coupling enables credibility-dependent actions (Schelling-style commitments, conviction-enabled persuasion, identity-enabled coordination); sustained over-coupling forecloses reality-dependent actions (the fully-captured cult member or motivated-reasoning subject loses access to actions that require accurate world-perception). At some interior point, $\mathcal{A}(\kappa_t)$ achieves its maximum — coupling has gone far enough to enable commitment-dependent actions, but not so far that reality-tracking is foreclosed — and beyond that point further coupling subtracts from the feasible set.

Discussion.

The empirical shape of $\mathcal{A}(\kappa_t)$ is plausibly **inverted-U**: increasing over $[0, \kappa^*]$ (commitment-enabling range; the agent gains credibility-dependent actions); decreasing over $[\kappa^*, 1]$ (over-coupling range; the agent loses reality-dependent actions). The inflection κ^* depends on the action class — for actions whose value depends primarily on credibility, κ^* is closer to 1; for actions whose value depends on accurate reality-tracking, κ^* is closer to 0. The shape is *open* — the framework does not currently derive its functional form, and this segment honestly flags the non-monotonicity as a structural commitment without committing to its quantitative shape.

The non-monotonicity has substantive implications: rational self-coupling is *bounded*. An agent that couples without bound forfeits the reality-tracking the commitment was meant to advance the agent's interests *in*. Schelling's commitment devices preserve flexibility along axes orthogonal to the commitment's target; effective intertemporal personal rules constrain a specific behavior while leaving general adaptive capacity intact; identity coupling that becomes total is the cult-formation pathology in which the agent loses access to non-identity-conforming action. The cluster's discipline — across all four mechanisms — is *partial* coupling along strategically-chosen axes, not maximal coupling everywhere.

Reversibility cost as the source of credibility.

The (M3) asymmetry is not incidental to strategic self-coupling — it *is* the operational core. Schelling's foundational treatment is exactly this point: the strategic value of a commitment derives from the costliness of its reversal. The same structural observation runs through Ainslie's personal rules (rules whose violation is psychologically expensive because of the precedent-breaking cost across the population of future-selves), Akerlof-Kranton's identity coupling (defections from identity-norms incur the identity-cost of becoming “the kind of person who would do this”), and Frank's emotional commitment (the ex-post irrationality of emotional response is its ex-ante credibility — a calculating agent would defect under sufficient incentive; the emotional agent cannot match that calculating defection because the emotion overrides the calculation).

Discussion.

The recurring structural shape across the four mechanisms: the agent couples to something *less fungible* than the original preference state. Schelling's external mechanism is less fungible than verbal commitment (the mechanism cannot be talked out of executing); Ainslie's personal rule is less fungible than per-occasion choice (breaking the rule has cross-temporal precedent cost); Akerlof-Kranton's identity is less fungible than incentive-driven preferences (identity-norms carry self-recognition costs incentive-violation does not); Frank's emotion is less fungible than calculation (the emotional response is structurally compelled, not chosen per-instance). The strategic move is *exchange of fungibility for credibility* — give up the per-instance optimization choice in exchange for the credible-commitment value that per-instance optimization cannot produce.

This frames a useful diagnostic. A claimed commitment whose reversal would be *cheap* is not a commitment; an interlocutor with appropriate skepticism will discount it. The agent who wants to acquire commitment-credible actions must accept the structural cost of binding to something less fungible than their current preference state. The mechanism choice (external structure / intertemporal rule / identity / emotion) determines *what kind* of fungibility is given up; the amount of fungibility given up determines the credibility acquired.

Asymmetric advantage in repeated strategic interaction.

Mirror to #disc-adversarial-coupling-pressure's asymmetric-advantage section: in repeated interaction, when two agents face each other, each has structural incentive to (a) acquire credibility-dependent actions through strategic self-coupling and (b) deny the opponent access to those same actions by undermining their commitments. The self-coupled agent acquires actions that the uncoupled opponent cannot match; the uncoupled opponent faces commitments that incentive-based negotiation cannot dissolve. If the asymmetry is sustained, the self-coupled agent compounds strategic reach — its commitments are read as credible; its threats and promises bind; its negotiating position is bounded below by its commitment, not above by its current incentive-state.

Discussion.

The asymmetric-advantage shape is dual to the adversarial-pressure case: there, the more-modular agent compounds informational advantage (cleaner observations of the less-modular opponent); here, the more-self-coupled agent compounds *strategic* advantage (credible threats and promises the uncoupled opponent cannot match). The two asymmetries can coexist — an agent might be defensively modular against adversarial information attacks *and* offensively self-coupled in narrow strategic axes (the diplomat who is genuinely informed about reality but credibly committed to specific positions; the negotiator whose information-processing is clear but whose walk-away point is structurally fixed). The architectural states are not contradictory — different axes of $\kappa_{\text{processing}}$ can carry different operations.

This is *structural* asymmetry rather than tactical: the operations on modularity state are not single-instance moves but *durable architectural commitments* that shape long-run interaction. The agent who has self-coupled to specific commitments is *that agent* for the duration of the coupling; reversing requires paying the (M3) uncoupling cost; sustained interaction therefore reads the commitments as fixed and adapts strategy accordingly. The self-coupled agent's strategic position is fundamentally different from the uncoupled agent's, in a way that incentive-based negotiation cannot capture.

Boundary with adversarial coupling pressure.

The boundary between strategic self-coupling (self-driven coupling-increasing) and adversarial coupling pressure (externally-driven coupling-increasing) blurs in practice. The cult-formation literature documents the canonical case: what begins as adversarial pressure (an external party manipulating the target's environment, social ties, information access) can be internalized as identity-coupling, after which the target perceives the coupling as self-driven and the line between operations is no longer cleanly drawn at the agent's experience. The behavioral signature of an agent at high $\kappa_{\text{processing}}$ may

not distinguish the two; the *structural* distinction (driver: self vs external) is recoverable only with information beyond the agent's behavior.

Discussion.

The framework's posture on the boundary: the *operations* are structurally distinct (driver is part of the operation's identity); their behavioral consequences can converge; the boundary is empirical when only behavioral signature is available; the boundary is structural when driver-history is observable. Cult-formation, social engineering at the identity-modification level, and propaganda-induced identity-binding all sit at the boundary; sustained successful pressure produces an agent who would internally classify the resulting coupling as self-chosen. The framework's machinery does *not* attempt to resolve the boundary at the level of agent experience — it characterizes the two operations as structurally distinct and admits that behavioral convergence is honest scope.

The relationship to truthification.

Strategic self-coupling is structurally the *dual* of truthification at a different layer of the agent. Truthification (instantiated via #der-class-coercion-via-wrapping and the institutional-scaffolding construction in #disc-adversarial-coupling-pressure §"Defensive scaffolding") is structural commitment to *goal-blind belief-update* — the W_1 strict wrapper commits the composite to operate Class 1 (Separated) at the belief-update layer by structurally foreclosing the goal-conditioned-update pathway. Strategic self-coupling is structural commitment to *goal-conditioned action-selection* — the commitment-device move commits the agent to operate in a particular action subspace by structurally foreclosing the deliberation-can-revise-this pathway.

Discussion.

Both are *wrapping constructions in the broader sense*. Both work by structural commitment that gives up flexibility along one axis in exchange for a different value along another. The two operations are duals: truthification gives up belief-pliability to gain reality-tracking; strategic self-coupling gives up action-pliability to gain commitment-credibility. The directed-separation property is foreclosed differently by each — truthification *restores* the separation at the belief-update layer; strategic self-coupling *deliberately weakens* it at the action-selection layer (specifically, the action-selection now depends on a coupled state κ_t that constrains $\mathcal{A}$, which is structurally a goal-conditioned narrowing of the action choice).

This dual relationship is structurally illuminating: both operations are deliberate departures from full architectural flexibility, in opposite directions on the goal-belief separation axis. The framework's epistemic-honesty discipline runs through both — truthification commits to *not letting goals corrupt belief*; strategic self-coupling commits to *letting goals constrain action* in stable, credible ways.

Epistemic Status.

Discussion-grade at the structural-recognition level. The segment articulates an extension of the framework's machinery (coupling-dependent action space; strategy-DAG enabling edges; reversibility-cost asymmetry) and adopts prior-art concepts (Schelling commitment devices, Ainslie personal rules, Akerlof-Kranton identity coupling, Frank emotions-as-commitment) as first-class theory components. The substantive claims are inherited from the prior-art literature at their original tiers:

- The credibility paradox (commitment value = uncoupling cost) is *exact* in Schelling's game-theoretic formalization for the stated mechanism-design setting; it is *robust qualitative* at AAT's level of abstraction (no AAT-internal derivation; the structural claim transfers across the four mechanisms).
- The inverted-U shape of $\mathcal{A}(\kappa_t)$ is *empirical / discussion-grade* — argued structurally, instantiated across the prior-art literatures, not derived from AAT primitives. Promotion would require a derived condition on the inflection point as a function of action-class composition.

- The reversibility-cost asymmetry is *robust qualitative* — observed across all four mechanisms, structurally derivable from the credibility-paradox observation.
- The dual relationship with truthification is *discussion-grade* — structural recognition that both operations are wrapping constructions in opposite directions on the goal-belief separation axis; the formal duality has not been derived.

Max attainable: *robust qualitative* for the structural pattern across the four mechanisms (the cross-mechanism convergence on the credibility-paradox shape is the kind of structural-recognition claim that promotes to robust-qualitative under careful prior-art search); *discussion-grade* for the inverted-U shape claim (promotion to *empirical* requires named domain-specific instantiation, e.g., the bounded-rationality literature’s empirical work on self-control reversals); *conditional-derived* for the dual-with-truthification recognition only if the formal duality is worked out.

Discussion.

The mirror-image relationship with #disc-adversarial-coupling-pressure. The two segments cover opposite-valence operations on the same architectural state. Adversarial pressure: externally-driven, vulnerability-creating, coupling expanding attack surface. Strategic self-coupling: self-driven, action-enabling, coupling expanding the agent’s own offensive reach. The two segments’ structural parallelism is deliberate — same architectural state ($\kappa_{\text{processing}}$ per #der-directed-separation), same operation direction (increasing κ), opposite driver (external vs self) and opposite strategic valence (vulnerability vs enabling). The framework’s third operation — truthification — is the opposite-direction operation (decreasing κ), instantiated via #der-class-coercion-via-wrapping (component-level wrapping) and #disc-adversarial-coupling-pressure §”Defensive scaffolding” (composite-level scaffolding). Together the three operations form the three-operation modularity-state-dynamics picture canonicalized at #disc-modularity-state-dynamics.

Connection to #norm-explicit-strategy-condition. The segment’s central operation — choosing to self-couple — is a strategic decision the agent makes about its own architecture. This connects to the explicit-strategy normative condition: when does the agent gain from deliberate strategy formation vs reactive exploration? Strategic self-coupling adds a new dimension to that question: when does the agent gain from deliberate *self-modification of architecture* vs leaving the architecture in its default state? The cost-benefit threshold for explicit strategy (#norm-explicit-strategy-condition) needs an analog for the self-coupling case — when does the credibility value of acquired actions exceed the structural cost of the coupling (the lost flexibility, the (M3) uncoupling cost amortized over the expected coupling duration)? An honest framing of the question is open; the segment names the structural dependency without deriving the threshold.

Composition consistency under self-coupling. Self-coupling within an agent may be analyzed as a composition move via #disc-composition-consistency. Ainslie’s intertemporal-bargaining framing makes this explicit: the agent treated as a population of temporally-displaced sub-agents who bargain; the personal rule is a coupling between current-self and future-self that fixes the future-self’s behavior in advance. The composite-agent reading: the agent at time t couples to its future-self at $t + 1$ by committing now to constrain $t + 1$ behavior. Formally this is a composite-agent move where the sub-agents are temporal slices of the same physical agent and the coupling is the personal-rule commitment. The connection is real but the formalization is open — the temporal-composite reading is suggestive but does not derive without explicit machinery for composite-agents-across-time. Worth following if the segment’s downstream uses make this connection load-bearing.

The action-space-as-coupling-state-dependent formalization $\mathcal{A}(\kappa_t)$ as the load-bearing structural extension. Of the three structural moves (M1)–(M3), the (M1) extension is the most consequential for downstream framework reach. Once the framework admits coupling-dependent action spaces, a range of phenomena currently outside the framework’s analytical surface become legible: commitment-device strategy; identity-driven coordination; emotion-as-strategic-resource; ritual and oath-taking as structural-commitment operations; institutional design that creates coupling-dependent action sets for collective actors. (M2) and (M3) are the supporting machinery. A future cycle that derives concrete

results from (M1) — e.g., a sector-condition analog for coupling-dependent action sets, or a persistence-condition variant for agents with non-stationary $\mathcal{A}(\kappa_t)$ — would promote this segment from discussion-grade to derived-conditional. The structural extension is the recognition; downstream derivation is the natural follow-on.

Findings.

Strategic Self-Coupling as the Enabling Polarity of Coupling-as-Property. **Brief:** When you make a promise that everyone can see, when you burn the bridge behind your army, when you adopt an identity that makes certain choices feel like betrayal of who you are, when you cultivate the kind of anger that *will* retaliate even when retaliation costs you — you are deliberately giving up flexibility. The thing you give up has a strategic value the flexibility could not produce. A threat that you could costlessly withdraw is not a threat; an opponent will ignore it. A promise you could renegotiate at any moment is not a promise; a partner will not bet on it. Across very different domains — Schelling’s commitment devices in international strategy, Ainslie’s personal rules in self-control, Akerlof and Kranton’s identity coupling in economics, Frank’s emotions as commitment devices — the same structural shape recurs: agents acquire credibility-dependent actions by coupling to something less fungible than their current preference-state, and the strategic value of the action they acquire is precisely the cost of reversing the coupling. The framework’s static picture of an agent’s modularity as architectural property is incomplete; modularity is a *contested, cultivated, and constructed* property that agents act on for their own strategic reasons. Adversaries press it down for exploitation (#disc-adversarial-coupling-pressure); agents press it down themselves for offensive reach (this segment); defensive scaffolding presses it back up (the truthification operation). All three are operations on the same state-space; what they do to it differs in driver and direction.

Impact: Names the second of three operations in the modularity-state-dynamics three-operation picture (cf. #disc-modularity-state-dynamics). Surfaces the framework’s previously-implicit fixed- $\mathcal{A}$ assumption (flagged in #disc-adversarial-coupling-pressure §”Scope”) and characterizes the structural extension required to handle coupling-dependent action spaces — the (M1)–(M3) extensions (coupling-dependent action space; strategy-DAG enabling edges; reversibility-cost asymmetry). Recognizes the mirror-image relationship with #disc-adversarial-coupling-pressure: same architectural state ($\kappa_{\text{processing}}$), same operation direction (increasing κ), opposite driver and opposite strategic valence. Identifies the *dual relationship* with truthification — both are wrapping-construction-style structural commitments giving up flexibility along one axis for value along another, but operating on opposite poles of the goal-belief separation. The dual recognition is structurally illuminating: the framework’s epistemic-honesty discipline runs through both — truthification commits to not letting goals corrupt belief; strategic self-coupling commits to letting goals constrain action in stable, credible ways. Opens the framework’s reach to a class of phenomena (commitment-device strategy, identity-driven coordination, emotion-as-strategic-resource, ritual and oath-taking, institutional commitment design) that the fixed-action-space framing did not admit.

Novelty Claim: *Claim recognition* of strategic self-coupling as a structurally-distinct operation on modularity state, sister to and opposite-valence with adversarial coupling pressure (#disc-adversarial-coupling-pressure); *claim differentiation* on the (M1)–(M3) structural extensions required to formalize the operation in AAT’s machinery (coupling-dependent action space, strategy-DAG enabling edges, reversibility-cost asymmetry); *claim transfer* of established commitment-device and identity-economics machinery into AAT’s analytical surface as first-class theory components. The empirical phenomena (Schelling commitment, Ainslie personal rules, identity coupling, emotional commitment) are well-documented in the cited literatures; the contribution is naming their place in AAT’s scope architecture, the structural extensions required, and the mirror-image relationship with adversarial coupling pressure that situates both operations within the three-operation modularity-state-dynamics taxonomy.

Related Work:

Table 11.4

<i>ASF concern</i>	<i>Prior-art language</i>	<i>Relationship / Positioning</i>
Commitment devices (external structure)	Schelling, T. C. (1960), <i>The Strategy of Conflict</i> , Harvard University Press; (1966), <i>Arms and Influence</i> , Yale University Press (found pre-2026)	<i>formal antecedent</i> — adopted with citation; the canonical commitment-device framework; the credibility paradox (commitment value = uncoupling cost) is Schelling’s foundational structural observation. Primary-source verification scheduled for promotion past draft
Intertemporal bargaining and willpower	Ainslie, G. (1992), <i>Picoeconomics: The Strategic Interaction of Successive Motivational States Within the Person</i> , Cambridge University Press; (2001), <i>Breakdown of Will</i> , Cambridge University Press (found pre-2026)	<i>formal antecedent</i> — adopted with citation; the agent-as-population-of-temporally-displaced-sub-agents framing; personal rules as the population-level coupling that stabilizes preference. Closer to AAT’s composition-consistency framing than Schelling’s external-facing version
Identity economics	Akerlof, G. A. & Kranton, R. E. (2000), “Economics and Identity,” <i>Quarterly Journal of Economics</i> 115(3):715–753; (2010), <i>Identity Economics: How Our Identities Shape Our Work, Wages, and Well-Being</i> , Princeton University Press (found pre-2026)	<i>formal antecedent</i> — adopted with citation; identity-utility-term as the formal hook for identity-coupled preferences; in-group coordination and costly signaling as the action class identity coupling enables
Emotions as commitment	Frank, R. H. (1988), <i>Passions Within Reason: The Strategic Role of the Emotions</i> , W. W. Norton (found pre-2026)	<i>formal antecedent</i> — adopted with citation; the ex-post-irrationality-equals-ex-ante-credibility structural observation; complement to Schelling for the <i>internal</i> -affective coupling case
Self-binding taxonomy	Elster, J. (1979), <i>Ulysses and the Sirens: Studies in Rationality and Irrationality</i> , Cambridge University Press; (2000), <i>Ulysses Unbound: Studies in Rationality, Precommitment, and Constraints</i> , Cambridge University Press (found pre-2026)	<i>adjacent literature</i> — survey treatment cataloguing forms of precommitment across psychology, politics, and law; useful for taxonomic comparison but not directly adopted
Evolution of commitment in social contracts	Skyrms, B. (1996), <i>Evolution of the Social Contract</i> , Cambridge University Press; (2004), <i>The Stag Hunt and the Evolution of Social Structure</i> , Cambridge University Press (found pre-2026)	<i>adjacent literature</i> — evolutionary game-theoretic treatment of commitment supporting cooperation; useful for population-dynamics layer when self-coupling becomes a population-level phenomenon (norms, institutions); Part III population-scope work could engage
Sister segment — vulnerability polarity	#disc-adversarial-coupling-pressure (this framework, 2026-05)	<i>sibling internal segment</i> — opposite-valence operation on the same architectural state; the structural-parallelism is deliberate; both operations are part of M4’s three-operation modularity-state-dynamics taxonomy at #disc-modularity-state-dynamics
Truthification operation — dual relationship	#der-class-coercion-via-wrapping (this framework); #disc-adversarial-coupling-pressure §”Defensive scaffolding” (this framework)	<i>internal antecedent</i> — the third M4 operation, structurally dual at the goal-belief separation axis (truthification commits to goal-blind belief-update; this segment’s operation commits to goal-conditioned action-selection)

Search Log:

- 2026-05-24 (*targeted, segment-grade*): Prior-art adoption work per the spike (spike-strategic-self-coupling.md §2.1 direct adoption candidates + §2.2 adjacent literatures). Primary citations imported at segment-grade — one-line characterizations plus full bibliographic citations; primary-source verification scheduled when the segment advances past draft. The structural recognition (four mechanisms instantiating one cross-domain pattern) is named here for the first time at the AAT-internal level; targeted future search candidates: (a) whether any prior synthesis names commitment-device / identity coupling / willpower / emotions-as-commitment as instances of one structural operation under formal nomenclature; (b) whether the action-space-as-coupling-state-dependent formalization $\mathcal{A}(\kappa)$ has been written down in active-inference or bounded-rationality literatures; (c) whether the dual-with-truthification framing has appeared in any framework engagement with the goal-belief separation question. Expected outcome: the constituent moves are well-precedented at the textbook level; the cross-mechanism unification as an operation on modularity state alongside its adversarial-pressure dual appears to be a fresh presentational recognition.
- 2026-05-09 (*intuition-only on the cross-domain unification*): The structural recognition that the four mechanisms (Schelling commitment, Ainslie personal rules, Akerlof-Kranton identity coupling, Frank emotional commitment) instantiate one structural operation at different coupling targets is named here. The spike (spike-strategic-self-coupling.md) established the recognition; this segment lands it canonically with the cross-domain pattern documented.

Working Notes.

- **Provenance.** Authored 2026-05-24 as Move 3 of the M4 modularity-as-contested-property cycle (msc/modularity-cycle-plan-2026-05-09.md §3 Move 3; PRACTICA priority #2). The spike at spikes/spike-strategic-self-coupling.md (2026-05-09) established the (P1) primary-segment scope, prior-art landscape, and (M1)–(M3) structural extensions; this segment lands the (P1) product. Move 4 — the M4 meta-segment #disc-modularity-state-dynamics — lands in the cycle's same commit cluster.
- **Open before candidate.** (a) Primary-source verification of the four canonical citations (Schelling 1960/1966, Ainslie 1992/2001, Akerlof-Kranton 2000/2010, Frank 1988). The segment imports at segment-grade depth; direct primary-source reading is queued. (b) The non-monotonicity / inverted-U shape of $\mathcal{A}(\kappa_t)$ is presented as a structural observation; a derived condition on the inflection point κ^* as a function of action-class composition would promote the claim from *discussion-grade* to *empirical* or *conditional-derived*. Open. (c) The reversibility-cost asymmetry's relationship to #der-deliberation-cost is named as adjacent but not formally derived; a derivation showing the asymmetry as a special case of deliberation-cost asymmetry under structural-commitment conditions would tighten the connection. (d) The composition-consistency reading of Ainslie's intertemporal bargaining (#disc-composition-consistency applied to temporally-displaced sub-agents) is suggestive but unformalized.
- **Cluster context.** This segment is the second-leg of M4 (#disc-modularity-state-dynamics). Move 4 — landing the M4 meta-segment itself — depends on this segment being in place. The meta-segment integrates this segment + #disc-adversarial-coupling-pressure (the first leg, landed earlier in the cycle) + the truthification mechanism via #der-class-coercion-via-wrapping and #disc-adversarial-coupling-pressure §"Defensive scaffolding".
- **Possible appendix companions** (per the spike §5 candidates P2 and P4). (P2) def-coupling-dependent-action-space — formalization appendix carrying $\mathcal{A}(\kappa_t)$ definition + monotonicity range + relationship to architecture classification + worked examples from prior art. (P4) Reversibility-cost-asymmetry formalization — only useful if it grounds a substantive derived result. Both deferred until the discussion-grade landing here gains downstream load-bearing uses.
- **Indexed at spike-proposals.** A sub-spike candidate flagged from this segment for spikes/PROPOSED.md consideration: the *inverted-U shape derivation* — deriving the inflection point κ^* from action-class composition (the fraction of credibility-dependent vs reality-dependent actions in the agent's relevant action space) as a function whose form predicts the agent's rationally-chosen self-coupling level. Open framework-extension question; appropriate as a Tier 2 / Tier 3 row when M4 lands and surfaces this as a follow-on.

- **Boundary with cult-formation literature.** The §”Boundary with adversarial coupling pressure” section flags that the line between self-driven and adversarial-driven coupling blurs at the agent’s level of experience (cult-formation; identity-conversion under sustained adversarial pressure). The framework’s posture is that the *operations* are structurally distinct (driver: self vs external) but their behavioral consequences can converge. Targeted follow-on work for that boundary lives in cult-formation literature (Lifton 1961 *Thought Reform and the Psychology of Totalism*; Singer 1995 *Cults in Our Midst*) which the segment cross-references at adjacent-literature depth.
- **Migration note (2026-05-09 GUC rename):** $\kappa_{\text{processing}}$ vocabulary is unchanged across the rename (the broader vocabulary lives at the operation level, not the Class label). The segment uses Class 1 (Separated) / Class 2 (Partial) / Class 3 (Coupled) in current canon vocabulary throughout. Removed at candidate stage per FORMAT.md Gate 4.

Incidental audit gold (lift 2026-05-30). Cross-audit “wandering thoughts” / §14-ideation harvested from the de-novo auditors’ working dirs, kept separate from the certified theory-fix findings. **Coverage:** as a late-authored Meta-Architecture II segment, only the deepest de-novo walk (773921, Gemini) reached it under its current slug — thin and single-substrate by expectation. Substrate attribution from the audit dir.

1. Candidate Brief prose / pre-prose.

- **Why perfect rationality is a handicap** — the auditor’s one-line framing of the segment’s payload: “a purely modular, perfectly rational, uncoupled agent is actually strategically handicapped in multi-agent environments because it cannot make credible promises or threats” (Gemini, AUDIT-WORKING-773921). A candidate Brief opener that inverts the naive “coupling is bad” reading.
- **The credibility paradox in words:** “a commitment’s value equals its uncoupling cost” — the prose gloss of the strict inequality $c_{\text{couple}} \ll c_{\text{uncouple}}$ (Gemini, AUDIT-WORKING-773921).

2. Candidate Discussion.

- **Emotions as commitment technology** (*strongest framing candidate here*). “Defining an emotion (like anger or love) as ‘coupling between rational calculation and affective response that trades fungibility for credibility’ mathematically explains why humans have emotions: an angry person who will retaliate regardless of cost is a more credible deterrent than a perfectly rational calculator who will accept a bribe” (Gemini, AUDIT-WORKING-773921, on the Frank 1988 leg). A candidate Discussion sharpening — the segment already adopts Frank; this is a pedagogically vivid restatement of *why* the affective-commitment leg matters.
- **The inverted-U as the framework finding its own limits.** The auditor read the $\mathcal{A}(\kappa_t)$ inverted-U (too much coupling eventually destroys reality-dependent actions $\rightarrow$ “cult-like delusion”) as “an excellent example of the framework finding the natural limits of its own operations” (Gemini, AUDIT-WORKING-773921). Candidate framing for the non-monotonicity Discussion.

3. Follow-up items.

- **Self-control as single-agent composition** — the auditor amplified the segment’s own Working-Note (d): treat an individual over time as a *composite agent* of temporally-displaced sub-agents (Self-at- t_1 , Self-at- t_2), so a personal rule (“I will not smoke”) is an inter-agent coupling agreement and AAT’s multi-agent composition math (11) applies directly to single-agent self-control (Gemini, AUDIT-WORKING-773921). Convergent confirmation that Working-Note (d)’s Ainslie / composition-consistency reading is worth formalizing.

4. Readers often ask / wonder.

- (No distinct reader-question surfaced beyond the items above in the reached audit.)

5. Candidate figures.

- (No dedicated figure proposed for this segment in the reached audit.)

Belongs elsewhere.

- **A unified vocabulary for institutional design.** Forward-vision: “contracts, constitutions, and cultural norms are all just technologies for manipulating $\kappa_{\text{processing}}$ to unlock new regions of the action space,” and the segment “mathematically formalizes why perfect rationality is evolutionarily unfit for social environments” (Gemini, AUDIT-WORKING-773921). Aspirational reach for the M4 framing layer (11.1) / field-positioning, not a per-segment edit.

11.3 Discussion: Adversarial Coupling Pressure

In adversarial settings, opponents do not merely deceive through the observation channel; they apply strategic pressure that drives a target *across architectural classes* — from Class 1 (Separated) toward Class 3 (Coupled). Coupling is itself an attack surface: a Class 1 (Separated) agent can be deceived only through the observation channel, while a GUC-Coupled agent exposes additional input ports through O_t and Σ_t . The result is a use-case expansion for 03-11m-core’s coupled formulation: the population that needs coupled-formulation analysis is broader than the architecturally-coupled. Any agent under sustained adversarial pressure that can drive its $\kappa_{\text{processing}}$ upward sits in the same scope, regardless of native class.

The phenomenon.

A Class 1 (Separated) agent’s f_M accepts inputs only through the event stream; its O_t and Σ_t have no legitimate input port from outside that stream. An adversary’s only attack channel is engineered observation ambiguity, and the bias they can induce is bounded by #deriv-observation-ambiguity-bias-bound’s constant C relative to the ambiguity they can manufacture.

Discussion.

A GUC-Coupled agent — Class 3 (Coupled), or Class 2 (Partial) with active $O_t \rightarrow M_t$ pathways — exposes additional channels. Three operational mechanisms recur across the literatures on propaganda, social engineering, marketing, and influence operations:

Table 11.5

<i>Mechanism</i>	<i>Channel exploited</i>	<i>What the adversary does</i>
Identity-binding	$O_t \rightarrow M_t$	Make the target <i>want</i> the false belief to be true (bind belief to identity, group membership, self-image). Revising M_t now requires revising O_t , raising the cost of correction.
Affect / urgency	bypass deliberate f_M	Manufacture time pressure or affective load that prevents the orient cascade from running cleanly; the agent acts before M_t updates resolve.
Sunk-cost engineering	$\Sigma_t \rightarrow M_t$	Induce the target to commit publicly or expensively to a strategy; subsequent M_t revisions that would invalidate Σ_t now carry a coupled cost (admit prior commitment was wrong), and the coupling bends M_t revision instead of Σ_t revision.

Each mechanism corresponds to a specific structural pathway through which adversary-controlled signals act on M_t via O_t or Σ_t — not through the legitimate observation channel.

Diagnostic corruption is directional.

The satisfaction-gap / control-regret split is silently corrupted under adversarial coupling pressure, and the corruption is *directional*: a competent adversary biases the diagnostic predictably rather than randomly.

Derived (Conditional on directional adversary objective; from #def-satisfaction-gap, #def-control-regret + adversarial-pressure mechanisms above).

If the adversary's interest is the target's continued effort along an unproductive axis, they bias the reading toward control regret — “the goal is right, try harder.” The target compounds effort where it cannot pay, while the adversary extracts whatever value the effort produces (attention, emotional engagement, financial commitment, political mobilization).

If the adversary's interest is the target's *withdrawal* from an axis they could win, they bias the reading toward satisfaction gap — “the goal was wrong, give up.” The target abandons achievable wins; the adversary captures the unclaimed ground.

Manufactured-outrage cycles instantiate the first; engineered-helplessness and demoralization campaigns instantiate the second. Both are *diagnostic capture*, not just deception about world-state — the target is reading a corrupted diagnostic and confidently acting on it.

Cascade inversion.

The orient cascade resolves observation $\rightarrow M_t \rightarrow \Sigma_t \rightarrow O_t$ in that order, with O_t revising only when forced. Adversarial coupling pressure inverts the cascade: the adversary acts on O_t first (manufacture desire / identity-stake / urgency), which forces Σ_t to align, which then makes legitimate observations flow *against* the gradient and get rejected as inconsistent with the now-committed Σ_t . The cascade still runs, but its information dependency is reversed — and in that reversed form, observation contradicting the pre-loaded objective is the *least* legitimate input the agent will accept.

Derived (Conditional on adversary controls O_t -targeted signal injection; from #der-orient-cascade).

Population expansion for coupled formulation.

The Class 3 (Coupled) formulation handed off to 03-llm-core/ was originally scoped to agents whose native architecture is non-modular: LLMs whose attention processes goals and observations together; biological cognition with motivated-reasoning pathways. Adversarial coupling pressure expands this population: a Class 1 (Separated) organization without internal $O_t \rightarrow M_t$ pathways but operating in a hostile information environment that targets its leadership's identity and commitments is, for analytical purposes, in the coupled-agent population during that engagement. The native architecture is irrelevant once the coupling exists in operation.

This is a *use-case expansion*, not a derivation expansion: the coupled-formulation machinery does not need new content for this case. What's new is the recognition that the population is broader than architecture-determined.

Scope: implicit assumptions and the polarity of coupling.

This segment covers one polarity of coupling-as-property — the *vulnerability* polarity, in which decreased modularity is externally driven and produces attack surface. Two implicit assumptions in the framework's machinery back the scope; both are worth surfacing rather than leaving hidden.

Discussion.

Bounded signaling. The framework formally restricts the $G_t \rightarrow$ world channel to action selection $a_t = \pi(M_t, G_t)$. A coarse $\mathcal{A}$ implicitly assumes π is sufficiently many-to-one — or the agent's behavioral signature sufficiently bandwidth-restricted — that observers cannot recover G_t from observable behavior alone. In practice the agent's full behavioral signature (micro-expressions, prosody, attention patterns, vocabulary choice, hormonal signals, response latencies) leaks

G_t at much higher bandwidth than the discrete-action channel suggests. Against capable adversarial inferers the bounded-signaling assumption fails, and this is part of *why* the bias bound's C becomes adversarially-tight in the §"Discussion" reframing: sophisticated adversaries effectively read G_t and target identity-bindings precisely. The assumption sits formally in #der-directed-separation's agent-world interface; it is not yet surfaced there as a load-bearing scope statement, and surfacing it back-pressures that segment as a follow-on item.

Fixed action set independent of coupling. The action space $\mathcal{A}$ in the policy $\pi(M_t, G_t) \rightarrow a_t$ does not depend on the agent's coupling state. This implicitly assumes the actions available to the agent are independent of $\kappa_{\text{processing}}$. In practice, some actions exist *only* under coupling: Schelling-style commitment devices (where credibility requires foreclosed retreat); the salesperson whose sincere conviction enables persuasion that the modular-but-skeptical salesperson cannot match; the marathoner who can run the race only because retreat options have been burned; religious practice requiring coupled belief-and-action; identity-economics models where coupled identity unlocks coordination outcomes. Coupling is not just a *cost*; it is also *enabling*. Treating $\mathcal{A}$ as fixed therefore mis-specifies the strategic landscape — it omits a class of actions that rational agents may *choose* to acquire by deliberately self-coupling, and it forces the framework to read all coupling as failure mode.

Three operations on modularity state. Together with the segment's core treatment of adversarial coupling pressure, the two assumptions surface a clean three-operation picture of modularity-state dynamics:

Table 11.6

<i>Operation</i>	<i>Driver</i>	<i>Effect on modularity</i>	<i>Effect on agent</i>
Truthification	self-driven	increases	restores diagnostic; opens Part II results; instantiated by <i>defensive scaffolding</i> below
Strategic self-coupling	self-driven	decreases	<i>enables</i> actions otherwise unavailable (credibility, conviction, commitment)
Adversarial coupling pressure	externally-driven	decreases	<i>creates vulnerability</i> (this segment's core treatment)

This segment's primary subject is the third operation. The first operation is instantiated in §"Defensive scaffolding as composition" below — institutional scaffolding as the truthification leg of the three. The second operation is currently outside the framework's machinery; see Working Notes for the candidate sister segment and the candidate three-operation meta-segment.

Defensive scaffolding as composition.

Restoring modularity under adversarial pressure is structurally a Part III move. Institutional scaffolding — peer review, prediction registers, double-entry bookkeeping, adversarial procedure, prediction markets, source-authentication discipline, structured red-teaming — constructs a composite agent (*individual + scaffolding*) with better Class 1 (Separated) properties than the bare individual, even when the individual's native architecture is Class 3 (Coupled) or driven there by adversarial pressure. The scaffolding externalizes channels of the orient cascade (belief-update checked by parties whose O_t doesn't share the individuals; prediction registers fixing M_t -claims before outcomes resolve; adversarial procedure forcing the feasibility check) so that resolution can run sequentially across the composite even when it cannot run sequentially within any single component.

Discussion.

The defensive composite recovers architectural modularity at the composite level whose properties exceed those of its sub-components — the routing-structure-driven side of the axis-decomposed composite-class-inheritance table in

#der-directed-separation Discussion: when the routing (the scaffolding's institutional protocol) is structured to externalize orient-cascade channels, a composite of Class 3 (Coupled) sub-agents (or sub-agents under adversarial pressure toward Class 3) acquires Class 1 (Separated) or Class 2 (Partial) properties at the composite level. This is the operational dual of the wrapping construction (**#der-class-coercion-via-wrapping**): wrapping operates at the *component* level by externally scaffolding a single component; defensive scaffolding operates at the *composite* level by structuring the routing across components. Both move the composite class along Axis A (architectural class per **#der-directed-separation**), independently of the dynamic-regime axis (**#disc-dynamic-regime-axis**) where goal-alignment effects live.

Epistemic Status.

Discussion-grade at the structural-recognition level. The segment articulates a use-case expansion of existing machinery rather than a derivation of new machinery; the substantive claims are inherited from upstream segments at their tiers:

- The Class 1 / 2 / 3 (Separated / Partial / Coupled) partition and $\kappa_{\text{processing}}$ measure are *robust qualitative* per **#der-directed-separation**.
- The bias bound that adversarial pressure attempts to saturate is *conditional* per **#deriv-observation-ambiguity-bias-bound** (under named sub-scopes).
- The diagnostic-corruption directionality is *robust qualitative*: it follows from **#def-satisfaction-gap** and **#def-control-regret** being arithmetic identities once the value objects are fixed, plus the structural observation that adversary preferences are typically directional rather than random; the directional bias is empirically observable across propaganda and influence-operation case studies.
- The orient-cascade inversion claim is *robust qualitative* — it follows from **#der-orient-cascade**'s information dependency under the explicit hypothesis that the adversary controls O_t -targeted signal injection.

Max attainable: *robust qualitative* for the structural pattern; the specific mechanism table is empirical (instantiated by propaganda, security, and influence-operation literatures) and would be promoted to *empirical* with named domain citations once a targeted prior-art search has been conducted.

Discussion.

This is the offensive complement to directed-separation-as-defense. **#der-directed-separation** characterizes modularity as architectural property. This segment names modularity as *contested* property under adversarial pressure: the architecture's modularity is what the adversary is trying to perforate, and what defensive scaffolding is trying to preserve or restore. The two segments together complete the picture — modularity is what the architecture enables (or fails to enable), and modularity is what the adversary attacks (or fails to attack).

Asymmetric advantage compounds in repeated interaction. When two agents face each other, each has structural incentive to (a) reduce its own coupling (defensive scaffolding) and (b) increase the opponent's coupling (offensive pressure). If the asymmetry is sustained, the more-modular agent gets clean observations of the less-modular agent's behavior — informative M_t updates that the bias bound's C does not impose — while the less-modular agent gets contaminated observations of the more-modular agent and bends M_t toward whatever pre-loaded objective the adversary engineered. The more-modular agent compounds informational advantage; the less-modular agent compounds bias. In repeated interaction with sustained asymmetry, this is *structural* rather than tactical advantage — modularity asymmetry produces strategic asymmetry.

Reframing the Class 3 (Coupled) bias bound's operational meaning. **#deriv-observation-ambiguity-bias-bound**'s constant C is presented in that segment as a worst-case bound. Adversarial coupling pressure reframes C as a *target* the adversary actively tries to saturate by engineering observation ambiguity to its maximum admissible value. Under sustained adversarial engineering, the bound is not slack — it is approached. This affects how the bound should be interpreted operationally: in non-adversarial settings, the bound is conservative; in adversarial settings, it is approximately

tight under competent adversarial design. Whether the bound *can* be saturated depends on what the adversary controls (the observation channel's ambiguity-engineering bandwidth); under unrestricted adversarial control the saturation is essentially complete.

Connection to the separability pattern. The architectural classification ladder of #disc-separability-pattern reads adversarial coupling pressure as an *operational* driver of population movement along the ladder: Class 1 (Separated) agents under sustained adversarial pressure functionally migrate toward the Class 3 (Coupled) / general-open tier for the duration of the engagement, even though their native architecture remains Class 1 (Separated). The structured-repair tier (Class 2, Partial) is then the natural home for *defensive scaffolding* — the explicit added machinery that restores identification under adversarial conditions. The separability pattern's three-tier shape is therefore not just a static catalog of architectural types; it is a phase-space across which adversarial pressure and defensive scaffolding move populations.

Findings.

Adversarial Coupling Pressure as Use-Case Expansion of Coupled Formulation. Brief: Adversaries do not just lie; they reshape the *architecture* of the agent they are attacking. A Separated agent (Class 1) — one whose beliefs update from observation alone — can only be deceived through what it sees; the adversary's reach is bounded. But an agent whose beliefs are bent by what it wants to be true exposes new attack surface: every channel through which the adversary can act on what the target wants, or on what plan the target has committed to, becomes a way to reach the target's beliefs without going through legitimate observation. Practical mechanisms — identity-binding, manufactured urgency, sunk-cost engineering — are operational instantiations of this single structural move: drive the target across the architectural boundary so the additional ports open up. The framework's machinery for Coupled agents (Class 3) (originally scoped to architectural cases like LLMs and motivated cognition) therefore covers a broader population than its initial framing suggested: any agent under sustained adversarial pressure, regardless of native architecture. Defensive scaffolding — peer review, prediction registers, double-entry bookkeeping, adversarial procedure — is the structural counter-move: it constructs a composite agent whose modularity properties exceed those of its individual components, restoring at the system level what cannot be sustained at the individual level.

Impact: Names a phenomenon currently treated only obliquely in AAT: the strategic incentive in adversarial interaction to drive opponent coupling. Reframes the Class 3 (Coupled) bias bound (#deriv-observation-ambiguity-bias-bound) from passive worst-case to adversarially-saturated target — under competent adversarial design the bound is approximately tight rather than conservative. Connects the orient cascade's clean sequential resolution to its adversarially-inverted form, where observation flows against the information gradient and contradictory inputs are the least legitimate ones the agent will accept. Identifies institutional scaffolding (peer review, prediction markets, double-entry bookkeeping, adversarial procedure) as a Part III composition move that restores Class 1 (Separated) properties at the composite level even when individual sub-components are coupled or driven coupled. Expands the natural population for 03-11m-core/s coupled formulation beyond architecture-determined cases to include any agent under sustained adversarial pressure.

Novelty Claim: *Claim recognition* of adversarial coupling pressure as a structural phenomenon in AAT's existing scope architecture — adversaries strategically drive coupling because coupling expands attack surface — combined with *claim differentiation* on the population scope of coupled-formulation analysis: not just architecturally-coupled agents, but any agent under sustained adversarial coupling pressure. The empirical phenomena (manufactured outrage, identity-binding, social engineering, sunk-cost-driven commitment) are well-documented in propaganda, security, and behavioral-economics literatures; the contribution is naming their place in AAT's scope architecture, the corresponding population expansion for coupled formulation, and the asymmetric-advantage-compounding implication for repeated adversarial interaction.

Related Work:

- *Likely adjacent literatures, not yet searched* — propaganda research (Bernays 1928 *Propaganda*; Ellul 1965 *Propagandes*; Stanley 2015 *How Propaganda Works*), social engineering (Mitnick & Simon 2002 *The Art of Deception*; Hadnagy social-engineering corpus), behavioral economics on motivated reasoning (Kunda 1990 *The Case for Motivated Reasoning*; Kahan and the cultural-cognition project), influence operations (Singer & Brooking 2018 *LikeWar*), and Bayesian persuasion / cheap talk (Kamenica-Gentzkow 2011; Crawford-Sobel 1982). Relationship: *adjacent literature, intuition-only* — these literatures document the empirical phenomena instantiating the mechanism table; the contribution this segment claims is structural placement and population expansion, not phenomenon discovery.

Search Log:

- 2026-05-09 (*intuition-only*): Empirical phenomena (manufactured outrage, identity-binding, social engineering, sunk-cost commitment dynamics) are well-documented across propaganda research, social engineering, behavioral economics, and influence operations. The AAT-internal placement (adversarial pressure as driver of $\kappa_{\text{processing}}$ saturation) and the population-expansion claim for coupled formulation have not been searched. Targeted search recommended before promotion to draft, focusing on: (a) whether adversarial-cognitive-architecture framing exists in the active-inference, control-as-inference, or game-theoretic-deception literatures; (b) whether the “adversary drives opponent coupling” framing has been explicitly named in the cognitive-warfare or information-operations academic literatures; (c) whether the asymmetric-advantage-compounding result has been derived in any repeated-game treatment of cognitive influence.

Working Notes.

- **Promotion: requires targeted prior-art search.** The structural-recognition claim is the segment’s contribution; the empirical phenomena are not novel. Until a Pillar-style targeted search is conducted, novelty status remains *intuition-only* and the segment should not advance past draft.
- **Stage $\times$ source $\times$ form structure within the adversarial mechanisms.** The mechanism table (identity-binding / affect-urgency / sunk-cost-engineering) is a projection of the broader Class 2 sub-typology of #disc-partial-coupling-pathways onto specific (stage, source) cells: identity-binding $\rightarrow$ P1+P2 / O / content-to-process (form varies with strength of identity-binding); sunk-cost $\rightarrow$ P3 / Σ / process (the cell that admits self-stabilizing belief-strategy attractors per #der-belief-strategy-attractor); affect/urgency $\rightarrow$ cascade-bypass (outside the static (stage $\times$ source $\times$ form) parameterization; a separate failure mode rooted in time-pressure-induced under-deliberation). The typology gives the structural reason these three mechanisms have distinct repair characters — and supplies the dynamical attractor structure that explains sunk-cost-engineering’s particular effectiveness in propaganda and influence operations (once the attractor is established, the target’s own machinery maintains the misaligned belief without further adversary input).
- **Brief is not yet at Feynman criterion.** The current Brief reaches for the structural recognition rather than an everyday physical analog whose causal structure is isomorphic to coupling-as-attack-surface. A candidate analog: the difference between selling someone a product (must convince them the product is good — bounded by observable quality) and selling someone the *desire* for the product (once they want it, they will not see the flaws — unbounded by observable quality). The salesperson who can drive desire-formation has acquired a channel into belief that does not pass through evidence about the product. Whether this analog carries the load-bearing structure (identity-binding specifically; the cascade inversion; the asymmetric-advantage compounding) needs work; the bathtub-equivalent has not been written.
- **Game-theoretic formalization candidate.** The “modularity asymmetry produces strategic asymmetry” claim in the Discussion is currently presented as a structural intuition rather than a derived result. A formalization candidate: define a two-agent repeated game where each agent’s payoff includes a term in the *other agent’s* $\kappa_{\text{processing}}$ (offensive value of opponent coupling) and a negative term in *own* $\kappa_{\text{processing}}$ (defensive cost of self coupling). Equilibrium analysis would predict (a) modularity arms races, (b) asymmetric advantage cascades when modularity asymmetry is sustained, and (c) defensive-scaffolding incentives.

The information-theoretic structure (the more-modular agent gets cleaner M_t updates from the same observation stream) is what drives the asymmetry. Plausibly tractable; would land as an appendix segment if pursued.

- **Defensive scaffolding deserves its own segment under Part III.** The institutional-scaffolding-as-composite-agent framing in §"Defensive scaffolding as composition" is a specific instantiation of #disc-composition-consistency with substantive content (peer review, prediction registers, double-entry bookkeeping, adversarial procedure, prediction markets, structured red-teaming all instantiate the same pattern: externalize a channel of the orient cascade so resolution runs sequentially across the composite even when it cannot run sequentially within any single component). Defer until the connection to existing composition results is mapped explicitly; a stub #deriv-defensive-scaffolding-composition may be the right home if the formalization tracks.
- **OUTLINE.md placement.** The segment is in 01-aat-core/src/ but is not yet in 01-aat-core/OUTLINE.md. Natural home: under Part II (it sits at the architecture-classification layer) or as a meta-segment alongside #disc-separability-pattern and #disc-identifiability-floor. The latter is probably the better fit — it's a meta-pattern that cuts across Part II and Part III rather than a Part II result proper.
- **Sister segment for strategic self-coupling.** The fixed- $\mathcal{A}$ assumption surfaced in §"Scope: implicit assumptions and the polarity of coupling" points at coupling's *enabling* polarity — actions that exist only under coupling (Schelling commitment devices, conviction-enabled persuasion, sunk-cost-foreclosed retreat, identity-coupled coordination outcomes). A sister segment covering this polarity directly is staged at spikes/spike-strategic-self-coupling.md. The natural prior-art adoption set is Schelling 1960 *The Strategy of Conflict* (commitment devices), Ainslie 1992 *Picoeconomics* (intertemporal bargaining and willpower), Akerlof-Kranton 2000/2010 (identity economics), Frank 1988 *Passions Within Reason* (emotions as commitment devices). Investigation pending; the spike establishes scope and prior-art landscape.
- **Meta-segment landed 2026-05-24.** The §"Scope" three-operation table (truthification / strategic self-coupling / adversarial coupling pressure) lands canonically at #disc-modularity-state-dynamics — the M4 meta-segment alongside M1 (#disc-identifiability-floor) / M2 (#disc-separability-pattern) / M3 (#disc-additive-coordinate-forcing). The three polarity-instantiations all have primary-segment homes now: this segment for adversarial coupling pressure; #disc-strategic-self-coupling for the enabling polarity; truthification via #der-class-coercion-via-wrapping (component-level) + §"Defensive scaffolding as composition" below (composite-level). The cycle landed the modularity-as-contested-property work per msc/modularity-cycle-plan-2026-05-09.md.
- **Surfacing bounded-signaling in #der-directed-separation.** The bounded-signaling assumption surfaced in §"Scope: implicit assumptions" is more naturally located in #der-directed-separation's Discussion as a load-bearing scope statement on the agent-world interface (it's about *that* segment's structure, not specific to adversarial pressure). Back-pressure follow-on: add to that segment a paragraph naming the assumption — that the channel from G_t to the world runs only through π , and that this assumption fails operationally when behavioral leakage is rich relative to action coarseness. Cross-reference back to this segment for the adversarial saturation case.
- **Migration note (2026-05-09 GUC rename + 2026-05-21 composite-claim withdrawal):** Class 2 $\leftrightarrow$ Class 3 swap. Pre-2026-05-09: Class 2 = fully merged, Class 3 = partially modular. Post: Class 2 = Partial, Class 3 = Coupled. This segment uses "coupling pressure" / "adversarial coupling" as its central concept — that broader vocabulary is unchanged. The specific GUC-class references were renamed: "toward Class 2 (fully merged)" $\rightarrow$ "toward Class 3 (Coupled)"; "Class-2 coupled formulation" $\rightarrow$ "Class 3 (Coupled) formulation"; "Class-3 architecture (partially modular)" for the defensive composite $\rightarrow$ architecturally-modular composite (routing-structure-driven per #der-directed-separation axis-decomposed inheritance); "Class-2 bias bound" $\rightarrow$ "Class 3 (Coupled) bias bound." The earlier "Class 1 (Separated) sub-agents $\rightarrow$ Class 2 (Partial) composite refinement" reference (inversion-comparator for the defensive scaffolding case) was withdrawn 2026-05-21 per spikes/strategic-composition-class-3-attempt-2026-05-21/ (Cournot witness against the unconditional claim; conflation diagnosis on belief-content vs processing-pathway); the defensive-composite passage in §"Defensive scaffolding as composition" was re-derived in regime-axis-aware language at the same time. Removed at candidate stage per FORMAT.md Gate 4.

Incidental audit gold (lift 2026-05-30). Cross-audit “wandering thoughts” / §14-ideation harvested from the de-novo auditors’ working dirs, kept separate from the certified theory-fix findings. **Coverage:** as a late-authored Meta-Architecture II segment, only the deepest de-novo walk (773921, Gemini) reached it under its current slug — thin and single-substrate by expectation. Substrate attribution from the audit dir. Several of the auditor’s strongest observations (defensive scaffolding as composition, the game-theoretic formalization, bounded-signaling) are *already* tracked in the Working Notes above; those are cross-referenced rather than re-stated.

1. Candidate Brief prose / pre-prose.

- **Coupling as attack surface** — the auditor’s compression of the segment’s frame: adversaries “don’t target what people *know* (the observation channel), but what people *want* or *identify with* (the $O_t \rightarrow M_t$ coupling channel)” (Gemini, AUDIT-WORKING-773921). A candidate Brief anchor (and a sharper handle than the current Brief’s structural-recognition phrasing — see the Feynman-criterion Working Note above).

2. Candidate Discussion.

- **Confirmation bias and demoralization as structural, not cognitive-quirk.** The auditor read “Cascade Inversion” as formally deriving these: by acting on O_t first (identity or manufactured urgency), the adversary “forces the Orient Cascade to run backward, making legitimate observations flow *against* the information gradient, guaranteeing their rejection” (Gemini, AUDIT-WORKING-773921). A candidate Discussion sharpening of the cascade-inversion mechanism — names the everyday phenomena (confirmation bias, demoralization) the structure explains.
- **Directional diagnostic corruption stated as two levers.** “If an adversary wants an agent to give up, they inflate the Satisfaction Gap (δ_{sat}); if they want it to waste effort on a doomed cause, they inflate Control Regret (δ_{regret})” (Gemini, AUDIT-WORKING-773921). A compact gloss linking the segment to 8.4 / 8.5.

3. Follow-up items.

- **Adversarial saturation of the bias bound.** The auditor stressed that the Coupled Bias Bound (AB) reads differently under adversarial conditions: “in a benign environment the bound C is a loose worst-case; in an adversarial environment the adversary is an optimizer actively trying to drive the agent’s error exactly to C ” (Gemini, AUDIT-WORKING-773921). A candidate sharpening of how the bias bound is invoked downstream — distinct from the static-bound reading. (The segment’s mechanism table already touches this; the “adversary-as-bound-optimizer” framing is the additional texture.)
- (Off-ramp note, not gold: the auditor’s other strong items — defensive scaffolding deserving its own composition segment, the game-theoretic two-agent κ -payoff formalization, and bounded-signaling belonging in 6.3 — are already tracked in the Working Notes above. No new follow-up needed.)

4. Readers often ask / wonder.

- **Are humans structurally weaker than text-only LLMs against adversarial coupling?** Following the bounded-signaling assumption: humans leak G_t via micro-expressions, prosody, and latency, bypassing the formal action channel a_t — “adversaries don’t have to guess what a human wants; they can read it directly,” whereas a text-only LLM leaks only via tokens (Gemini, AUDIT-WORKING-773921). A natural reader question that sharpens the bounded-signaling Working Note’s operational stakes.

5. Candidate figures.

- (No dedicated figure proposed for this segment in the reached audit.)

Belongs elsewhere.

- **A formal vocabulary for information warfare / propaganda / cult formation.** Forward-vision: the segment “extends AAT from a theory of agents into a theory of security against cognitive attack,” giving “a formal, structural vocabulary for Information Warfare, Propaganda, and Cult Formation based on the physics of the OODA loop” (Gemini, AUDIT-WORKING-773921). Aspirational reach for the M4 framing layer (11.1) / field-positioning, not a per-segment edit.
 - **Peer review and double-entry bookkeeping as Class-1-from-Class-3 composition** — the auditor independently praised the §”Defensive scaffolding as composition” framing (“not just administrative overhead; structural Part III composition moves that create a Class 1 macro-agent out of Class 3 human sub-agents”) (Gemini, AUDIT-WORKING-773921). Converges on the existing Working-Note proposal to give defensive scaffolding its own Part III segment; recorded as cross-substrate confirmation of that already-tracked move.
-

Chapter 12

Scope and Composition Formation

12.1 *Scope: Multi-Agent Scope*

Part III's broadest scope: any system of multiple agents each satisfying 1.6 and interacting through a shared environment. Each agent has its own model and purposeful substate; each observes the environment and other agents' actions; each acts; their actions jointly determine the next environment state. The general case for organizations, teams, ecosystems, and adversarial encounters. *Coupling goes through the shared environment, not through direct state modification* — agents affect each other by affecting the environment.

A methodological default the framework adopts: **correlated observations are generic; independence is the special case requiring justification**. When agents share an environment, their observations are generically correlated — they see aspects of the same reality. Independence (uncorrelated observations) requires the agents to observe non-overlapping aspects, which is the special case.

Each agent's observation decomposes into environmental and *inter-agent* components — messages from other agents. The **multi-agent routing structure** is a key new object: the *topology* (who communicates with whom) plus the *protocol* (the rule governing what class of information flows between agents). The protocol is a *rule specifying the channel*, not the content of any message — individual messages reflect senders' states (including individual goals) through their policies. What the routing structure governs is the *infrastructure* — which channels exist and what kind of information they carry. **Routing is goal-blind** when neither topology nor protocol depends on the composite's goal state. *Goal-dependent routing* (crisis-specific channels, mission-specific intelligence-sharing protocols, reorganizing reporting chains based on objective) breaks goal-blindness and is the lever that distinguishes Case 1 from Case 2 in 13.3.

The framework draws a *load-bearing scope distinction* between two classes of machinery. **Agent-level machinery** (individual persistence, agent tempo, per-agent mismatch, sector-condition stability) applies to *every* agent in *every* multi-agent configuration — cooperative, adversarial, or indifferent — because each agent individually satisfies 1.6. **Composite-level machinery** applies only when 12.2 is satisfied. The adversarial regime spans the negative-teleological-unity end of the unity spectrum and is treated *both* ways: equilibrium-convergent adversarial pairs are captured by composite-level machinery (the strategic-composite route in 16.1); cyclic or non-convergent adversarial dynamics are captured by agent-level machinery applied across the pair.

Formal Expression.

A multi-agent system consists of N agents $\{A_1, \dots, A_n\}$, each satisfying the scope condition (1.6), interacting through a shared environment with state $\Omega_t \in \mathcal{S}_{env}$:

Scope (multi-agent-scope).

- Each agent A_i has state $X_t^{(i)} = (M_t^{(i)}, G_t^{(i)})$
- Each agent observes: $o_t^{(i)} = h^{(i)}(\Omega_t, a_t^{(-i)}, \xi_t^{(i)})$ — observations may depend on other agents' actions
- Each agent acts: $a_t^{(i)} = \pi^{(i)}(X_t^{(i)})$
- The environment evolves: $\Omega_{t+1} = T(\Omega_t, a_t^{(1)}, \dots, a_t^{(n)}, \omega_t)$

The coupling is through the environment: agent i 's actions enter agent j 's observation function and the shared environment transition. Agents may also communicate directly (a special case of action-observation coupling with a dedicated channel).

Observation decomposition and routing.

Each agent's observation decomposes into environmental and inter-agent components:

Definition (observation-decomposition).

$$o_t^{(i)} = \left(o_{\text{env},t}^{(i)}, \{m_{ji,t}\}_{j \in \mathcal{N}_t(i)} \right) \quad (12.1)$$

where: - $o_{\text{env},t}^{(i)} = h_{\text{env}}^{(i)}(\Omega_t, \xi_t^{(i)})$: direct environmental observation (no inter-agent content) - $\mathcal{N}_t(i) \subseteq \{1, \dots, N\} \setminus \{i\}$: the **communication neighborhood** — which agents send messages to i at time t - $m_{ji,t} = c_t^{(j \rightarrow i)}(X_t^{(j)})$: message from j to i , determined by the sender's full state and the communication protocol

The **multi-agent routing structure** $R_t = (\mathcal{N}_t, \{c_t^{(j \rightarrow i)}\})$ specifies: - The **topology** $\mathcal{N}_t$: who communicates with whom
- The **protocol** $c_t^{(j \rightarrow i)}$: the rule governing what class of information flows from j to i

Definition (multi-agent-routing-structure).

Note: the protocol $c_t^{(j \rightarrow i)}$ is a *rule* specifying the channel, not the specific content of any message. Individual messages reflect the sender's state $X_t^{(j)}$ — including their individual goals — through the sender's policy. What the routing structure governs is the *infrastructure*: which channels exist and what kind of information they carry. *Bare-prose shorthand: the term "routing structure" is sanctioned within this segment after the first compound-form introduction; cross-segment citation should use the full "multi-agent routing structure" form.*

Routing is **goal-blind** when neither the topology nor the protocol depends on the composite's goal state:

Definition (goal-blind-routing).

$$\mathcal{N}_t \perp G_t^c \quad \text{and} \quad c_t^{(j \rightarrow i)} \perp G_t^c \quad \forall j, i \quad (12.2)$$

This means the communication infrastructure does not change based on what the composite is trying to achieve. Individual messages naturally reflect individual agents' goals through their policies — this is action, not routing. The routing condition is about the *structure* of information flow, not the *content* of individual messages.

Goal-dependent routing occurs when either the topology or the protocol varies with G_t^c . Examples: activating crisis-specific communication channels, changing intelligence-sharing protocols based on the current mission, reassigning reporting chains based on the operational objective.

Epistemic Status.

Axiomatic. This is a scope definition — it describes the class of systems Part III addresses. The only substantive choice is that coupling goes through the shared environment rather than through direct state modification. This follows from the agent boundary assumption (1.1): agents affect each other by affecting the environment, not by directly altering each other's internal states.

Discussion.

Correlated observations as default. When agents share an environment, their observations are generically correlated — they see aspects of the same reality. Independence (uncorrelated observations) requires the agents to observe non-overlapping aspects of the environment, which is the special case. Most multi-agent settings of interest involve substantial observation correlation.

The adversarial case sits along a spectrum — partly inside, partly outside the composite scope. Agents whose objectives conflict are multi-agent systems with negative teleological unity (14). Whether they form a composite depends on whether their strategic interaction admits an equilibrium structure: equilibrium-convergent adversarial pairs (potential or monotone games per Monderer-Shapley / Rosen) satisfy 12.2 via route (C-iv) as **strategic composites** with equilibrium-based macro-state — see 16.1 for the equilibrium-theoretic machinery and the (SC-1)–(SC-3) decomposition. Cyclic or non-convergent adversarial pairs satisfy none of (C-i)–(C-iv) and remain within 12.1 only. The asymmetric attacker / target case (one agent treated as an exogenous parameter rather than a sub-agent running its full AAT loop) is also a 12.1 phenomenon, handled by 15.2.

Two distinct classes of machinery apply across this spectrum:

- **Agent-level machinery** (individual persistence, agent tempo, per-agent mismatch, sector-condition stability) applies to *every* agent in *every* multi-agent configuration — cooperative, adversarial, or indifferent — because each agent individually satisfies 1.6 and Part I/II results apply directly. The adversarial tempo advantage (15.4) and adversarial destabilization (15.2) results are applications of this agent-level machinery to the case where one agent's actions are a disturbance source for another.
- **Composite-level machinery** applies *only* when 12.2 is satisfied via at least one of (C-i)–(C-iv). The (C-i)–(C-iii) alignment routes give composites with a shared objective O_c and admit closure-defect / team-persistence / composite-tempo / unity-closure machinery via 13.1. The (C-iv) strategic route gives composites with an equilibrium-based macro-state and admits the equilibrium-convergence / sector-template / closure-defect machinery in 16.1; the alignment-route and strategic-route forms are structurally distinct (an equilibrium-convergent adversarial composite has no shared O_c , only a shared $\mathcal{E}$).

The adversarial regime spans the negative- U_O end of the unity spectrum. Equilibrium-convergent adversarial pairs are captured by (C-iv) strategic-composite machinery; cyclic / non-convergent adversarial dynamics and asymmetric attacker / target configurations are captured by agent-level machinery applied across the pair. The existing adversarial segments (15.2, 15.4) operate at the agent-level for cases that fall outside (C-iv); 16.1 operates at the composite-level for cases that fall inside it.

Inter-agent communication as a special observation channel. Messages from j to i are actions by j and observations by i . The routing structure formalizes the infrastructure: who talks to whom ($\mathcal{N}_t$) and under what protocol ($c_t^{(j \rightarrow i)}$). The sender controls the content (unlike passive environmental observation), which introduces strategic manipulation (14.4). The gain from inter-agent communication enters the distributed tempo (14.4, Working Notes).

The routing/content distinction matters for directed separation. Individual messages reflect senders' goals — that's just action through policy. The directed-separation question at the composite level (13.3) is about the *routing structure*: does the infrastructure change based on the composite's goals? Goal-blind routing preserves directed separation; goal-dependent routing breaks it.

Working Notes.

Incidental audit gold (lift 2026-05-30). Cross-audit “wandering thoughts” / §14-ideation harvested from the de-novo auditors’ working dirs, deduplicated across substrates and lightly attributed by substrate + audit. Orthogonal pedagogical / framing / figure / naming material, kept separate from the certified theory-fix findings (handled elsewhere — see report flags for this segment). **Coverage:** four dirs reached a digested reflection on this segment (193847, 526815, 829314, 849201), plus the §III composition-foundations batch (471203). The other contributing dirs did not reach a dedicated reflection here or stopped earlier. Substrate attribution inferred from voice where not explicit; uncertain cases hedged.

1. Candidate Brief prose / pre-prose.

- The routing/content distinction stated as a hardware/software analogy: the routing structure $\mathcal{N}_t$ is “the organizational equivalent of separating network hardware from application software” — “a company can have a perfectly objective, goal-blind *email system* (routing) even if every individual *email* is highly biased by the sender’s personal goals” (Gemini, AUDIT-WORKING-193847). Independently praised as “masterful” / the segment’s standout move by three substrates (Gemini, AUDIT-WORKING-193847; Gemini, AUDIT-WORKING-829314 — “the mathematical key to understanding organizational epistemology”; Claude, AUDIT-WORKING-849201).
- A formal-taxonomy-of-groups gloss for an opener: three-tier ladder of “groups” — (1) independent agents (uncorrelated Ω); (2) multi-agent systems (correlated Ω , no coherence); (3) composites (correlated Ω , structural coherence) (Gemini, AUDIT-WORKING-193847).
- “You cannot apply composite-level theorems (like composite tempo or team persistence) to a bar brawl” — a vivid one-line scope guard for the multi-agent-vs-composite boundary (Gemini, AUDIT-WORKING-193847).

2. Candidate Discussion.

- **Goal-dependent routing as organizational pathology — the “wartime CEO” trade.** A worked elaboration of the Case-2 failure: a “wartime CEO” who restructures channels to serve the current objective (forms a Tiger Team, mandates “only report blockers related to Project X”) enslaves the information infrastructure to the goal — data about other failures is no longer routed to anyone, and the organization becomes a Class 3 (Coupled) agent suffering “institutional confirmation bias at the hardware level.” The framing: AAT gives the exact math for why organizational focus “trades long-term epistemic accuracy for short-term strategic tempo” (Gemini, AUDIT-WORKING-829314; the same CEO-fires-QA mechanism at Gemini, AUDIT-WORKING-193847 — severing the channel that “delivered aporia”). A candidate Discussion vignette making the goal-blind-vs-goal-dependent routing distinction operationally concrete. (*Note the early finding-vs-framing texture: the auditor reads this as a derived consequence; the segment treats it at hypothesis/discussion grade via 13.3 — keep the tier honest if promoted.*)
- **The calibrate-against-bias point.** Even when message *content* is biased (a department pads its numbers), a goal-blind *channel* lets the receiver “learn to calibrate against the bias” — separating the durable infrastructure question from the per-message-honesty question (Gemini, AUDIT-WORKING-829314). Sharpens why routing-blindness, not message-honesty, is the load-bearing condition.

3. Follow-up items.

- **Forward-reference / topological-sort stumble.** Two substrates independently flagged that the Discussion’s detailed treatment of routes (C-i)–(C-iv) reads as jarring out of order, since 12.2 (where those routes are defined) is the *next* segment — “it seems the author wrote or revised this discussion *after* writing the next segment” (Claude, AUDIT-WORKING-829314 suggested moving the route detail forward to 12.2 and leaving this segment purely as the un-fused multi-agent substrate definition). A reader-flow signal, not a defect; worth a sentence orienting the reader that the routes are defined next.
- The C-iv “strategic composite” framing (enemies in a stable equilibrium modeled as one macro-agent) was repeatedly flagged as the most counter-intuitive move and as “requiring very careful prose handling later” — see the richer note under 12.2.

4. Readers often ask / wonder.

- “If the Cold War is a single composite agent, what is its O_t ?” Multiple substrates converged on this as the natural reader question the C-iv route provokes: a strategic composite has no shared O_c , only a shared equilibrium $\mathcal{E}$ — “a fascinating extension of agency to non-teleological structures” (Gemini, AUDIT-WORKING-193847; Gemini, AUDIT-WORKING-829314; Claude, AUDIT-WORKING-849201). A candidate readers-often-ask preemption (the answer already lives in the (C-iv) macro-state-relative-to- $\mathcal{E}$ formalism — surface it where the reader first meets the strategic-composite idea).

5. Candidate figures.

- **Two-layer routing diagram.** Agents + shared environment at the bottom; routing infrastructure $\mathcal{N}_t$ as an overlay on top. The visual point: message *content* may be goal-colored while the routing *layer* can still be goal-blind or goal-dependent — the figure separates infrastructure from behavior (Codex/Claude, AUDIT-WORKING-526815, “diagram attempt”; converges with the hardware/software analogy above).

Belongs elsewhere.

- **Consciousness-infrastructure imperative: guaranteed goal-blind routing for a society-of-mind.** For Zi-am-tur or any multi-agent consciousness infrastructure, the infrastructure *must* guarantee goal-blind routing: sub-agents (safety monitor, planner, memory retriever) need communication channels that cannot be severed just because their messages are inconvenient — “if the infrastructure allows the planner to dynamically sever channels to silence dissent, the composite intelligence will inevitably collapse into epistemic closure” (Gemini, AUDIT-WORKING-193847). Aspirational reach pointing at 03-11m-core / 04-eli-core / architecture (and resonant with 13.3’s OPSEC observation), not at this scope segment.

12.2 Scope: Composite Agent

The scope condition that distinguishes a *composite agent* from a *multi-agent system*. A set of purposeful sub-agents, each satisfying 1.6, constitutes a *composite agent* only when their objectives exhibit sufficient teleological alignment to define a coherent composite purpose. Without this condition, the sub-agents form a multi-agent system (12.1) that may still be analyzed — but the *composition machinery* (13.1, 15.1, 13.2) does not apply because there is no composite agent whose quantities to compute. This parallels the single-agent 1.6: before asking how an agent adapts, one must check that an agent is *present*.

The scope condition is satisfied via *any of four disjunctive routes*. **(C-i) Shared composite objective** — the strongest alignment: all sub-agents’ policies are approximately optimal under a common objective. **(C-ii) Hierarchical derivation:** sub-objectives are derived by structure-preserving decomposition from a parent objective (military chain of command; corporate department structure) — sub-agents may not individually know the parent. **(C-iii) Mutual-benefit alignment** — the weakest route that still qualifies: a relevance variable exists such that joint actions raise each sub-agent’s expected value above the non-cooperation baseline (symbiotic coexistence; commensal ecologies; trading partners). **(C-iv) Equilibrium-convergent strategic interaction:** an equilibrium concept (Nash, correlated, coarse correlated) exists such that coupled best-response dynamics converge to its support — qualitatively distinct from the first three because it requires *neither* shared objectives *nor* hierarchical derivation *nor* mutual benefit, only structural convergence of the strategic interaction. Composites satisfying (C-iv) are **strategic composites**; their macro-state is defined relative to the equilibrium structure rather than to a shared target. The routes are *not* shown to reduce to a common scalar threshold; each carries its own operationalization, and only one is required for scope-satisfaction.

A specific load-bearing role: the framework’s third identifiability-floor instance (5.1, Instance 3 — composite contraction certification from component-marginal data is structurally impossible) has four structural escapes (observable coupling topology; matched Tier; passivity certificate; common contraction metric), but each escape needs a coherent target to certify. **Scope-satisfaction via at least one of the four routes is the enabling condition** for any of these escapes: without it, the “composite” is ill-defined (O_c does not exist under (C-i)-(C-iii); equilibrium structure $\mathcal{E}$ does not exist under (C-iv)) and the escapes have no coherent target.

Formal Expression.

Given sub-agents $\{A_1, \dots, A_N\}$ each satisfying 1.6, they constitute a *composite agent* iff there exists a common objective structure under at least one of the following routes:

Scope (scope-composite-agent).

(C-i) Shared composite objective. There exists O_c such that each sub-agent's effective policy is ϵ -compatible with O_c -optimal:

$$\exists O_c : \forall i, D(\pi_i, \pi_i^{O_c}) \leq \epsilon \quad (12.3)$$

where $\pi_i^{O_c}$ is sub-agent i 's optimal policy under the shared objective and D is an appropriate policy divergence. Strongest route; the sub-agents are all optimizing for the same thing, possibly with local decompositions.

(C-ii) Hierarchical derivation. There exists a parent objective O_c from which sub-objectives $\{O_i\}$ are derivable by decomposition consistent with 11:

$$\{O_i\} = \mathcal{D}(O_c) \quad (12.4)$$

where $\mathcal{D}$ is a structure-preserving decomposition. Each sub-agent optimizes its own O_i , but all O_i trace back to O_c . Military chain of command; corporate department structure. Sub-agents may not individually know O_c .

(C-iii) Mutual-benefit alignment. There exists a relevance variable Y such that the sub-agents' joint actions raise $\mathbb{E}[Y]$ above the non-cooperation baseline for each sub-agent:

$$\exists Y : \forall i, \mathbb{E}[Y \mid \text{joint}] > \mathbb{E}[Y \mid \text{non-coop}] \quad (12.5)$$

Weakest route. No explicit common objective, but interactions are positive-sum in some dimension. Symbiotic coexistence; commensal ecologies; trading partners who share no goals beyond mutual benefit.

(C-iv) Equilibrium-convergent strategic interaction. There exists an equilibrium concept $\mathcal{E}$ (Nash, correlated, or coarse correlated) such that coupled best-response dynamics of $\{A_i\}$ converge to (or cycle within the support of) $\mathcal{E}$:

$$\exists \mathcal{E} : \text{coupled best-response dynamics converge to the support of } \mathcal{E}. \quad (12.6)$$

Qualitatively distinct from (C-i)–(C-iii): requires neither shared objectives, nor hierarchical derivation, nor mutual benefit. Requires only **structural convergence** of the strategic interaction in the game-theoretic sense. The sub-agents' objectives may be partially opposing, but the interaction admits a stable joint behaviour. Composites satisfying (C-iv) are **strategic composites**, distinguished from alignment composites (C-i, C-ii) and mutual-benefit composites (C-iii). The composite's macro-state is defined relative to the equilibrium structure $\mathcal{E}$ rather than relative to a shared target state. See #deriv-strategic-composition for the A2'-analog transfer of #result-sector-persistence-template under potential-game (Monderer-Shapley 1996) or monotone-game (Rosen 1965) conditions, the (SC-1)–(SC-3) existence/stability/convergence decomposition, and the honest sub-scope β' scope exit for non-potential non-monotone games (VI existence + regret-minimization CCE set-convergence only).

Disjunctive form.

The scope condition is satisfied when **any** of (C-i), (C-ii), (C-iii), or (C-iv) applies. (C-i)–(C-iii) are progressively weaker qualitative requirements for what “teleological alignment sufficient to define a coherent composite purpose” means; (C-iv) covers strategic interaction with partially-opposing objectives via equilibrium convergence rather than alignment. The routes are *not* shown to reduce to a common scalar threshold. Each route carries its own operationalization and its own N -agent aggregation:

Scope (scope-composite-agent, disjunctive).

- (C-i) uses a value-function divergence $D(\pi_i, \pi_i^{O_c})$ aggregated across sub-agents.
- (C-ii) uses decomposition consistency of a parent objective O_c — a structural check, not a scalar.
- (C-iii) uses the existence of a relevance variable on which each sub-agent’s marginal contribution is positive — a per-pair existential check, not a magnitude.
- (C-iv) uses existence of an equilibrium structure $\mathcal{E}$ under coupled best-response dynamics — a fixed-point check (Nash / VI / regret-minimization CCE), structurally distinct from the alignment checks in (C-i)–(C-iii).

The teleological unity measure U_O from 14 (pairwise value-correlation aggregated to the group) tracks one projection of alignment — primarily route (C-i) — but is not a reduction of all three routes to a single scalar. Downstream segments (14.1, 12.3, 15.1) describe quality *conditional on scope-satisfaction* without assuming a common threshold: they presume the scope condition holds via at least one route, then analyze composite quantities within that regime.

(C-i) gives the strongest alignment; (C-iii) gives the weakest that still qualifies. A composite may satisfy multiple routes simultaneously; only one is required for scope.

What fails the scope condition: sub-agents with orthogonal objectives that also fail to admit equilibrium convergence — no shared or derivable O_c , no relevance variable providing mutual benefit, and no equilibrium concept (pure Nash, mixed Nash, or CCE) whose support the coupled best-response or no-regret dynamics reach — or unclassifiable objective-structure coupling. Such systems remain within 12.1 but not 12.2. Adversarial pairs that admit equilibrium convergence via (C-iv) — whether pure-strategy Nash under α' (potential / monotone games), mixed Nash under β' (Nash 1950 existence for finite games), or CCE in distribution under β' (Hart-Mas-Colell 2000) — DO satisfy composition-scope-condition as strategic composites. Cyclic games (rock-paper-scissors, matching pennies) lack a pure-strategy Nash but admit mixed Nash and CCE convergence; they fall within β' and satisfy (C-iv). The narrow category that fails (C-iv) is games with no equilibrium concept whose support is reachable by any admissible dynamic — a genuinely small class within the standard game-theoretic landscape.

Epistemic Status.

Robust-qualitative. Max attainable: *robust-qualitative.* The composite-agent scope is a structural proposal paralleling the single-agent scopes (1.5 and 1.6, both axiomatic — the minimal conditions for adaptive and purposeful machinery to be non-vacuous). Unlike them, 12.2 is a theoretical choice: we could admit teleologically misaligned groups as composites (via the closure-defect framework alone) and separately track whether their composite objective is meaningful. The proposal here is that this separate tracking is better done as a scope restriction: composition quantities are defined for systems satisfying at least one of (C-i)–(C-iv), not for systems satisfying none.

The three alternative routes are not exhaustive but cover the well-understood cases. The relationship between the routes (whether they partition or overlap) and whether there is a single underlying scalar that reduces them all is open; the disjunctive form above is chosen because no such reduction has been demonstrated.

Discussion.

Why this is a scope condition, not merely a quality metric. The existing 13.1 framework treats any group admitting a low- ϵ^* projection as a composite. This is mechanically defensible but runs into a category problem: a composite's purposeful substate $G_c = (O_c, \Sigma_c)$ requires a well-defined composite objective O_c . If the sub-agents' objectives do not align under any common O_c , then G_c is ill-defined and the composite is a fiction. Closure defect can be computed (as a projection property), but calling the result a “composite agent” strains the term. Making this a scope condition resolves the category issue: composition applies where G_c is well-defined.

Relationship to the single-agent scopes. The parallel is tight. The single-agent scopes (1.5 + 1.6) together say: before asking whether an agent persists, check that it is an agent (it observes under residual uncertainty, and at least one action has causal effect). 12.2 says: before asking whether a composite persists, check that it is a composite (teleological alignment sufficient to define O_c). In both cases, the scope precedes the quantitative machinery.

Relationship to 14. U_O in 14 has a dual role: (i) scope tracking — primarily the (C-i) route, where value-correlation is the operationalization; (ii) quality parameter for action-closure ϵ_a , via 14.1. The two roles are not in tension, but they are also not identical across routes: scope-satisfaction via (C-ii) or (C-iii) does not directly correspond to a U_O value, so U_O is best read as a quality metric *conditional on scope*, not as the defining scope variable.

Four routes: three alignment routes plus one strategic-equilibrium route. Routes (C-i), (C-ii), (C-iii) — the **alignment routes** — differ in how directly the sub-agents share purpose. They are progressively weaker qualitative requirements — (C-i) strongest, (C-iii) weakest — but they are not shown to reduce to a single scalar. (C-i) uses value-correlation directly; (C-ii) uses hierarchical decomposition consistency; (C-iii) uses marginal benefit existence. A composite can satisfy multiple alignment routes simultaneously (a well-aligned team satisfies both (C-i) and (C-ii)). Route (C-iv) — the **strategic-equilibrium route** — is qualitatively distinct: it admits composites with partially-opposing objectives whose strategic interaction is equilibrium-convergent (potential or monotone games per Monderer-Shapley / Rosen; see #deriv-strategic-composition), and its composite macro-state is defined relative to the equilibrium structure $\mathcal{E}$ rather than a shared target. What matters for scope is that at least one of the four routes holds.

Miller's IAM definition. Miller (2022, *Ex Machina*, Ch 1) proposes that social behavior requires Interaction, Agency, and Mutual benefit (IAM). The mutual-benefit requirement in (C-iii) is the AAT analog of Miller's third element. Miller treats IAM as the definition of social behavior; the proposal here treats alignment as the scope condition for composite-agent analysis. The two align substantively: Miller's mutual benefit is a minimum teleological unity.

Multi-agent systems with and without composite status. Adversarial pairs partition along a finer line than alignment / non-alignment. Adversarial pairs that admit equilibrium convergence — via pure-strategy Nash under α' (potential or monotone strategic games per #deriv-strategic-composition), mixed Nash under β' (Nash 1950 existence for finite games), or CCE in distribution under β' (Hart-Mas-Colell 2000 under no-regret dynamics) — satisfy the scope condition via (C-iv) as strategic composites. Cyclic games (rock-paper-scissors, matching pennies) sit inside β' : they lack pure-strategy Nash but admit mixed Nash and CCE convergence, so they satisfy (C-iv). The narrow category that genuinely fails (C-iv) is games with no equilibrium concept whose support is reachable by any admissible dynamic — a small class within the standard game-theoretic landscape. The asymmetric case — one agent as exogenous attacker, one as target — is a #scope-multi-agent phenomenon handled by #der-adversarial-destabilization; it does not form a composite in either (C-i)–(C-iii) or (C-iv) senses because the attacker is treated as a parameter rather than a sub-agent running its full AAT loop. Symmetric strategic composition under partially-opposing objectives is covered by (C-iv) and forms a strategic composite with equilibrium-based macro-state.

Load-bearing enabling role under #disc-identifiability-floor Instance 3. The composition-layer identifiability-floor instance establishes a no-go: composite contraction $\kappa_c > 0$ is not in general certifiable from component-level data alone — the coupling-sign bit distinguishing cooperative from adversarial regimes is unidentifiable from component marginals. Four structural escapes are available (observable coupling topology; matched Tier at composite level; passivity certificate; common contraction metric). Scope-satisfaction via at least one of (C-i)–(C-iv) in this segment is the **enabling condition** for any of these escapes: without scope-satisfaction, the “composite” is ill-defined (O_c does

not exist under (C-i)–(C-iii); equilibrium structure $\mathcal{E}$ does not exist under (C-iv); R_c is ill-typed), and the escapes have no coherent target to certify. The scope condition is therefore load-bearing apparatus that the composition-layer floor depends on — scope-satisfaction is what distinguishes a composite whose contraction or equilibrium convergence can be certified (under some escape) from a multi-agent system whose “composite contraction” is not a well-defined object.

How composites come into being: symbiogenesis. The formal mechanism by which a group passes from failing all three routes to satisfying at least one is described in 12.3. Briefly: in biological and social symbiogenesis, a host integrates an endosymbiont; the endosymbiont’s objective is subsumed into the host’s; the combined entity’s objective O_c emerges where two independent objectives stood. After the transition, (C-ii) applies: the endosymbiont’s objective is a sub-objective of O_c . This is the dynamical process of composite-agent identity creation; the scope condition here describes its result.

Working Notes.

- **Migration note (2026-05-09 GUC rename):** Class 2 $\leftrightarrow$ Class 3 swap. Pre-2026-05-09: Class 2 = fully merged, Class 3 = partially modular. Post: Class 2 = Partial, Class 3 = Coupled. Removed at candidate stage per FORMAT.md Gate 4.
- **Whether a common scalar exists across the alignment routes (resolved 2026-05-21 — negatively).** The current form treats the three alignment routes (C-i)–(C-iii) as a disjunction of qualitatively distinct conditions rather than as instances of a unified threshold; the strategic-equilibrium route (C-iv) is structurally distinct from all three. The 2026-05-21 strategic-composition spike at spikes/strategic-composition-class-3-attempt-2026-05-21/ resolved this open question in the *negative*: the four scope routes correspond to *regime-tier entry conditions on the dynamic-regime axis* per #disc-dynamic-regime-axis, not to a continuous scalar threshold. Specifically: (C-i)–(C-iii) enter R0 contraction-regime; (C-iv) with potential/monotone game-structure enters R1 equilibrium-regime; (C-iv) without potential/monotone structure enters R2 cyclic-distributional-regime; population-scope limit enters R3 mean-field-equilibrium-regime. The disjunction is structurally appropriate because the regimes themselves are structurally distinct (different fixed-point structures, different Lyapunov-certificate types, different macro-state types per the #disc-dynamic-regime-axis §”Identity” structural-identity for finite N). A unifying scalar across all four would mean a continuous reduction of categorically-distinct regime classes — which the regime-axis treatment shows is not the right organizing structure. U_O remains a quality metric for (C-i) within R0, not a unifying threshold across the routes.
- **Asymmetric unity.** Routes (C-i) and (C-ii) implicitly assume symmetric alignment (sub-agents equally represented in O_c). Symbiogenesis, firm acquisitions, and hierarchical delegation all feature *asymmetric* unity (host dominates; endosymbiont subordinate). The scope condition should accommodate asymmetric alignment; the current (C-ii) hierarchical form partially covers it, but an explicit treatment of role-asymmetric composites is open.
- **Interaction with Class 3 (Coupled) architectures.** 6.3’s architectural classification (Class 1 Separated, Class 2 Partial, Class 3 Coupled) applies to individual agents. Composites of Class 3 (Coupled) agents — e.g., multi-LLM systems — may satisfy this scope condition via shared training rather than shared explicit objective. Whether (C-i)–(C-iii) cover this adequately, or whether a fourth route for implicit alignment is needed, is open.
- **Transitivity.** If A_1, A_2 align and A_2, A_3 align, do A_1, A_3 ? Probably yes via (C-ii) if there is a common parent; not necessarily via (C-i) without a shared objective. Transitivity would make composite formation more robust; non-transitivity would explain why coalition-building requires explicit alignment work.
- **Promotion dependencies.** Before claims-verified: the common-scalar question (first bullet) is now resolved negatively (the scope routes correspond to dynamic-regime entry conditions, not to a unifying scalar — see #disc-dynamic-regime-axis); the remaining dependency is to audit whether any existing Part III segment implicitly assumes this scope without naming it (likely yes for team-persistence, tempo-composition, closure-based results) and update each to reference 12.2 for scope-satisfaction explicitly, including whether they apply to (C-i)–(C-iii), to (C-iv), or to both. Per the regime-axis treatment, (C-i)–(C-iii) $\rightarrow$ R0 / (C-iv) $\rightarrow$ R1 or R2 / population-scope $\rightarrow$ R3 — segments can now name their applicable regime directly.

Incidental audit gold (lift 2026-05-30). Cross-audit “wandering thoughts” / §14-ideation harvested from the de-novo auditors’ working dirs, deduplicated across substrates and lightly attributed. Orthogonal pedagogical / framing / figure / naming material, kept separate from the certified theory-fix findings (handled elsewhere — see report flags). **Coverage:** four dirs reached a digested reflection here (193847, 526815, 773921, 849201) plus the §III composition-foundations batch (471203); 829314 skipped this segment. This was a high-engagement segment — the four-route disjunction and the C-iv strategic-composite move drew rich ideation. Substrate attribution inferred from voice where not explicit.

1. Candidate Brief prose / pre-prose.

- **The “Why are we working together?” four-route taxonomy** — the standout plain-language gloss, converged on by two substrates: (C-i) shared objective = “because we want the same thing” (cult / startup); (C-ii) hierarchical derivation = “because the boss told us to” (corporation / military); (C-iii) mutual benefit = “because it’s profitable for both of us right now” (market transaction); (C-iv) strategic equilibrium = “because we are trapped in a Mexican standoff” (Cold War) (Gemini, AUDIT-WORKING-193847; the same four-domain mapping — team / corporation / trade network / stable ecosystem — at Claude, AUDIT-WORKING-773921 and Claude, AUDIT-WORKING-849201). A strong Brief-field seed for the segment.
- The scope condition as “the *legal boundary* for when we are mathematically permitted to zoom out and treat a group of agents as a single macro-agent” (Claude, AUDIT-WORKING-849201) — a compact framing of what the segment buys.
- The disjunctive-not-scalar form repeatedly praised as “an act of deep epistemic honesty” / “extreme mathematical maturity” — refusing to collapse qualitatively distinct organization types “into a single number” prevents the fallacy of composition (Gemini, AUDIT-WORKING-193847; Claude, AUDIT-WORKING-849201). Worth preserving as the anchor framing for *why* the disjunction is right.

2. Candidate Discussion.

- **C-iv as the formal boundary of “emergence.”** Stated sharply: “emergence is when a system satisfies (C-iv) and therefore can be modeled as a single macro-agent” — the scope condition converts an informal philosophical word into a checkable condition (Claude, AUDIT-WORKING-773921). A candidate Discussion sharpening of what C-iv contributes beyond game theory.
- **The non-teleological-composite point, made concrete.** “Do not assume a composite agent has a single, conscious mind or a single explicit objective O_c .” A free market of trading partners (C-iii) or a stable predator-prey ecosystem (C-iv) is formally a composite agent because its macro-dynamics are bounded and predictable, “even though no individual sub-agent holds the macro-objective” (Claude, AUDIT-WORKING-773921). Candidate Discussion guard against over-reading “agent.”
- **Symbiogenesis as the blueprint for safely building large composites** (aspirational, but with a structural spine). “You cannot safely build a monolithic super-intelligence from scratch (it hits the $\alpha > \rho/R$ wall because its R is too large to maintain high α). You must build it symbiotically” — start from a crèche of small bounded agents interacting via (C-iii) mutual benefit, then solidify into a (C-ii) hierarchy as protocols stabilize (Gemini, AUDIT-WORKING-193847). The structural claim (large R defeats high α , so composition is the route) is real; the application is reach — see Belongs-elsewhere. (*Avoid the “super-intelligence” vocabulary if promoted — canon prefers “future AI.”*)

3. Follow-up items.

- **Asymmetric / role-asymmetric composites under-modeled.** Symbiogenesis, firm acquisitions, and hierarchical delegation are *asymmetric* (host dictates O_c ; endosymbiont subordinate); the current (C-ii) only partially covers this, and explicit host-parasite power-dynamics modeling is “fertile ground for expansion” (Claude, AUDIT-WORKING-773921). Converges with the segment’s existing “Asymmetric unity” Working Note.
- **Transitivity failure as a structural explanation of coalition fragility.** If A_1 aligns with A_2 and A_2 with A_3 , alignment is non-transitive under the (C-iii) mutual-benefit route (a shared *parent* via (C-ii) would give transitivity; mutual benefit alone does not). “This mathematically explains why broad political coalitions are so fragile: they rely on non-transitive pairwise alignments” (Gemini, AUDIT-WORKING-193847). A candidate sharpening of the segment’s existing “Transitivity” Working Note into a stated consequence.

4. Readers often ask / wonder.

- “Does symbiogenesis happen via parametric learning or a structural snap?” A natural reader question once composition is defined — slow weight-alignment vs a sudden topological transition (Claude, AUDIT-WORKING-849201). The next segment (12.3) and its weighted-Lyapunov $\mu \rightarrow 0$ sketch begin to answer this; surface the pointer.
- “Is a strategic composite (e.g. the Cold War) really one agent?” Converged reader stumble on C-iv (see the parallel note under 12.1); the answer is the equilibrium-relative macro-state, worth preempting at first encounter.

5. Candidate figures.

- **Gate-into-four-routes diagram with a split output.** Multi-agent scope as a gate opening onto four routes, with the output deliberately split: C-i through C-iii yield an *objective-based* composite (shared O_c); C-iv yields an *equilibrium-based* strategic composite (shared $\mathcal{E}$, no O_c). The two target types kept visually separate because downstream machinery must know which object it computes over (Codex/Claude, AUDIT-WORKING-526815, “diagram attempt”).

Belongs elsewhere.

- **04-eli-core / consciousness-infrastructure — the crèche-to-hierarchy developmental path.** The symbiogenic-construction reading (“provide the stable environment where N independent agents can slowly fuse into one composite agent without collapsing into a non-convergent adversarial regime during the transition”) is a developmental blueprint pointing at the crèche / nursery material in 04-eli-core/, not at this scope segment (Gemini, AUDIT-WORKING-193847). Converges with the cost-shadow nursery reading lifted into 4.5.
- **Prior-art mapping — Miller (2022) IAM.** The (C-iii) mutual-benefit route maps onto Miller’s IAM (Interaction, Agency, Mutual benefit) social-behavior definition; flagged as bridging Agent-Based Modeling with Control Theory (Gemini, AUDIT-WORKING-193847; Claude, AUDIT-WORKING-849201). A prior-art/Discussion-citation pointer rather than gold per se — the segment already references Miller; this is a convergence signal that the mapping reads as natural to fresh readers.

12.3 Hypothesis: Symbiogenic Composition

Symbiogenesis is an asymmetric composition mechanism in which one agent (the *host*) integrates another (the *endosymbiont*) as a specialized sub-component, with the endosymbiont’s objective gradually subsumed into the host’s. It is distinct from peer coupling (13.1) and from population-level restructuring (the extreme transition motif drawn from Miller 2022, discussed in 4.8): symbiogenesis is how composite agents *come into existence* by crossing the 12.2 from below. The mechanism is well-attested empirically (eukaryote formation, firm mergers, legal-precedent adoption, language families) but formally underspecified within AAT. This segment captures the phenomenon and flags the specific formalization gaps.

Formal Expression.

Given two purposeful agents A_h (host) and A_e (endosymbiont), each satisfying 1.6, symbiogenic composition is a process on the joint state space with three coupled dynamics:

Hypothesis (symbiogenic-composition).

(S-1) Objective absorption. The endosymbiont’s objective O_e transforms toward alignment with or derivation from the host’s objective O_h :

$$O_e(\tau) \xrightarrow{\tau \rightarrow \tau_{\text{consolidated}}} \mathcal{D}_e(O_h) \quad (12.7)$$

where $\mathcal{D}_e$ is a derivation functional (in the sense of route (C-ii) in 12.2): O_e becomes a sub-objective derived from O_h . Before: O_h and O_e are independent objectives, no route of 12.2 applies, and the pair is a multi-agent system (12.1) rather than a composite. After: O_e is a role within O_h ; route (C-ii) applies; the composite (A_h, A_e) satisfies the composition scope condition.

(S-2) Function transfer. Structural content from the endosymbiont's state (elements of M_e or Σ_e) transfers to or becomes accessible by the host:

$$\{M_h, \Sigma_h\}(\tau) \xrightarrow{\tau \rightarrow \tau_{\text{consolidated}}} \{M_h, \Sigma_h\} \cup \mathcal{F}(M_e, \Sigma_e) \quad (12.8)$$

where $\mathcal{F}$ is a transfer mapping (structure-preserving integration of endosymbiont functions into host state). In biological symbiogenesis: gene transfer. In organizational symbiogenesis: acquired firm's processes, patents, know-how integrated into acquirer's operations. This is the grafting operation of 9.7 in its cross-agent form — the host grafts structure originating in the endosymbiont.

(S-3) Autonomy reduction. The endosymbiont's effective action space contracts; many of its choices become fixed by the host's coordination:

$$\mathcal{A}_e^{\text{effective}}(\tau) \xrightarrow{\tau \rightarrow \tau_{\text{consolidated}}} \mathcal{A}_e^{\text{restricted}} \subsetneq \mathcal{A}_e^{\text{initial}} \quad (12.9)$$

The endosymbiont retains enough autonomy to avoid catastrophic transfers (e.g., mitochondria retain some genome to handle local fast-timescale responses that would be hazardous to route through the host nucleus) but loses most independent decision-making.

Integrated transition. At consolidation, the joint system is a single composite agent A_c whose substate contains the integrated structure:

$$X_c = (M_c, G_c) = (M_h \cup \mathcal{F}(M_e, \Sigma_e), (O_c, \Sigma_c)) \quad \text{with } O_c \approx O_h \quad (12.10)$$

The endosymbiont persists as a specialized sub-component of the host, not as an independent agent. The 12.2 is now satisfied; the peer-coupling machinery of 13.1 applies to the resulting composite.

Epistemic Status.

Robust qualitative. Max attainable: *robust qualitative* — the phenomenon is well-attested empirically across biological and social domains, but a general mathematical formalization within AAT is open.

What is well-established (externally):

- The existence of symbiogenesis as a distinct evolutionary mechanism (Mereschkowsky 1905, 1910; Sagan 1967; Margulis & Sagan 1997). Mitochondria and chloroplasts are the paradigm cases.
- The social analog in firms, technology, language, legal systems, religions (Miller 2022, Appendix B).
- “Innovation by parts” as qualitatively different from “innovation by sparks” (gradual mutation) — Miller's framing.

What is *not* derived within AAT:

- A formal model of the objective-transfer dynamics (S-1). What evolutionary or optimization process drives $O_e \rightarrow \mathcal{D}_e(O_h)$?

- A formal specification of the transfer functional $\mathcal{F}$ in (S-2). What structure is preserved, what is lost, what is transformed?
- A precise characterization of autonomy reduction (S-3). Why does the endosymbiont retain some autonomy rather than becoming fully deterministic?
- Quantitative predictions — e.g., when symbiogenesis is favored over peer coupling, what governs the timescale of consolidation, under what conditions it reverses.

The three dynamics (S-1), (S-2), (S-3) are proposed schemas, not results. A follow-up development of an AAT-specific dynamical model is the natural next step.

Discussion.

The role of this mechanism in Part III. Three distinct composition mechanisms are now in scope:

1. **Peer coupling** (13.1, 15.1, 13.2) — sub-agents interact through shared environment; closure defect measures faithfulness of projection. Presumes scope-satisfaction via at least one route of 12.2 (not a scalar U_O threshold).
2. **Extreme transition motif** (Miller 2022; introduced in 4.8; pending dedicated segments for composition-transition dynamics, latent structural diversity, and endogenous coupling) — population-level restructuring via neutral drift / niche creation / cascading displacement. U_O shifts across a population as agent types replace one another.
3. **Symbiogenesis** (this segment) — hierarchical absorption. U_O crosses the composition scope condition from below, creating a composite that did not previously exist.

Before symbiogenesis, the sub-agents were a multi-agent system (12.1) but not a composite. After, the resulting composite is subject to all of AAT's composition machinery. The symbiogenic transition is the specific dynamical process of composite-agent identity creation.

Why this cannot be modeled as peer coupling. Peer coupling assumes pre-existing sub-agents being projected into a macro-description. The closure-defect framework presupposes the composite exists; it measures how faithfully the macro tracks the micro. Symbiogenesis is about a composite *coming into being* from two previously-independent agents. No projection Λ of the pre-symbiogenic system yields the post-symbiogenic composite, because the endosymbiont's objective *changes* during the process — its objective is different before and after. The transformation is intrinsic to the sub-agents' state, not an external projection choice.

Why this cannot be modeled as extreme transition. The extreme transition motif operates at the population level with many agents, neutral drift of types, and niche-construction dynamics. Symbiogenesis is typically between two specific agents (or a small number) and proceeds through a specific asymmetric integration rather than through statistical population dynamics. The mechanisms overlap — symbiogenesis often occurs as part of a larger transition — but the core mechanism of symbiogenesis (bilateral asymmetric integration) is distinct from the population-level dynamics of extreme transitions.

Examples across domains.

Table 12.1

<i>Domain</i>	<i>Host</i>	<i>Endosymbiont</i>	<i>Integrated composite</i>
Biology	Archaeal host cell	α -proteobacterium	Eukaryotic cell (mitochondrion persists as organelle)
Biology	Eukaryotic cell	Cyanobacterium	Plant cell (chloroplast persists as organelle)
Commerce	Acquiring firm	Acquired firm	Merged firm (acquired operates as division)

<i>Domain</i>	<i>Host</i>	<i>Endosymbiont</i>	<i>Integrated composite</i>
Technology	Base platform	Integrated component	Composite product (component operates within host system)
Linguistics	Host language	Adopted vocabulary/grammar	Creolized / evolved language
Law	Legal system	Adopted precedent	Evolved jurisprudence (precedent operates as doctrine)
Religion	Host tradition	Absorbed elements	Syncretic practice

In each case: asymmetric integration, autonomy reduction of the absorbed entity, gradual objective subsumption, functional specialization.

Connection to 9.7. The single-agent “grafting” operation in 9.7 is within-agent — an agent incorporates external structure into its own Σ_t . Symbiogenesis is cross-agent — the grafted structure originates in another agent, and the integration is accompanied by that other agent’s objective being absorbed. These are related but distinct: grafting is the structural-change mechanism on the host side; symbiogenesis is the bilateral process that includes grafting plus objective-absorption plus autonomy reduction.

Rate-distortion interpretation (connecting to 14.1). Under the Information Bottleneck conjecture in 14.1, peer coupling is IB compression with the relevance variable defined by a shared composite objective. Symbiogenesis is the process by which the relevance variable itself shifts: from two separate IB problems (each sub-agent’s own survival objective) to a single IB problem (the composite’s survival objective). The symbiogenic transition creates the shared relevance variable, which in turn makes the IB frontier well-defined for the composite. This is a structural shift in the IB problem, not a compression along a fixed IB frontier.

Working Notes.

- **Objective-transfer dynamics (S-1).** The most load-bearing open formalization. What process drives $O_e \rightarrow \mathcal{D}_e(O_h)$? Candidates: evolutionary selection (endosymbionts whose objectives align with host survival are selected for, since the alternative is extinction); bounded-rationality constraint (coordinating two divergent objectives exceeds the endosymbiont’s capacity, forcing simplification); explicit design (firm mergers where acquired objectives are deliberately restructured). Each gives a different dynamical equation.
- **Function transfer $\mathcal{F}$ (S-2).** Needs to respect the structure of the host’s M_h and Σ_h . In biology, gene transfer preserves molecular functions but changes regulatory context. In social analogs, the analog is: functions are preserved, but their triggers and dependencies change. A general specification is open.
- **Autonomy reduction (S-3).** Why not complete? The endosymbiont retains some autonomy because complete integration would eliminate the fast local response capacity that made symbiogenesis advantageous in the first place. A cost-benefit analysis on autonomy retention (in the style of 9.9) would make this quantitative.
- **(S-3) as weighted-Lyapunov limit (sketch-level).** J’s asymmetric limit $\alpha_2 \rightarrow 0$ under weighted Lyapunov $V_\mu(\xi) = \frac{1}{2}(\|\delta_1\|^2 + \mu\|\delta_2\|^2)$ with $\mu \rightarrow 0$ formalizes (S-3): the endosymbiont’s autonomous correction dynamics are weighted out of the joint stability argument, leaving the host’s sector condition as the composite’s persistence condition. This is a smooth deformation of the peer-coupling (CM4) inequality, not a discontinuous regime change — symbiogenesis and peer coupling are parameter-limits of the same weighted-Lyapunov analysis. The sketch is promotable to derived once (S-2) function transfer is formalized in this segment (the weighted Lyapunov limit does not address what happens to M_h when structure from M_e is inherited).
- **Reverse symbiogenesis.** Endosymbionts occasionally regain autonomy (biological examples: some organelle-hosted genes return to the nucleus; organizational examples: acquired divisions spun off). Theoretically: the scope condition can be crossed in either direction. A composite that loses U_O dissolves back into a multi-agent system. The triggering conditions and typical dynamics are open.

- **Interaction with logogenic agents.** In LLM-based agent architectures, multiple models can compose through shared training or through interface-specified protocols. Whether this constitutes symbiogenesis (with one model dominating) or peer coupling depends on whether the component models retain independent objectives. Worth investigating in 03-llm-core/.
- **Quantitative predictions.** When is symbiogenesis favored over peer coupling? Hypothesis: when the coordination overhead C_{coord} between would-be peer-coupled agents exceeds the integration cost of symbiogenesis. Transaction-cost theory (Coase / Williamson) is the economic analog. The AAT version would connect C_{coord} (13.2) to the energetic or informational cost of maintaining separate objectives, with symbiogenesis favored when the latter exceeds the former.
- **Timescale of consolidation.** In biology, symbiogenesis takes evolutionary time (millions of years). In firms, months to years. In software/ideas, potentially much faster. The consolidation timescale $\tau_{\text{consolidated}}$ is domain-dependent; a general characterization is open.
- **Saddle-node bifurcation analysis — quantitative threshold form, conditional on a nonlinear coordination penalty.** Under the *hypothesis* that aggregate multi-agent mismatch δ obeys $\dot{\delta} = \rho_{\text{env}} - \alpha_{\text{auto}}\delta + k\delta^2$ with a nonlinear coordination penalty $+k\delta^2$ ($k > 0$) capturing compounding coordination failure between misaligned autonomous sub-agents, the steady-state fixed points $\delta^* = (\alpha_{\text{auto}} \pm \sqrt{\alpha_{\text{auto}}^2 - 4k\rho_{\text{env}}})/(2k)$ exist only for $\rho_{\text{env}} \leq \rho_c := \alpha_{\text{auto}}^2/(4k)$. Above ρ_c the two fixed points collide in a *saddle-node bifurcation* and disappear: $\dot{\delta} > 0$ everywhere, and the autonomous multi-agent system has no stable equilibrium. The symbiogenic escape route is structural: the composite merges sub-agent objectives and state ($\mu \rightarrow 0$), eliminating the coordination-penalty term and recovering linear dynamics on the merged state. The bifurcation analysis therefore predicts symbiogenesis as a *mathematically forced phase transition* at critical environmental volatility ρ_c rather than as a contingent organizational choice. **Status: conditional on derivation of the $+k\delta^2$ coordination penalty from #def-shared-intent.** The $+k\delta^2$ form is currently stipulated rather than derived; making it rigorous requires showing that compounding coordination failure between agents with mismatch δ_A, δ_B produces an aggregate-mismatch dynamics with this specific quadratic-in- δ structure under named conditions on the shared-intent quantity. Without that derivation, the threshold $\rho_c = \alpha_{\text{auto}}^2/(4k)$ is a *formulation*, not a theorem; with it, the threshold form is exact under the named hypothesis. The closed-negative result here — that the bifurcation derivation is conditional on a derivation step not yet attempted — is itself load-bearing: any future strengthening of #hyp-symbiogenic-composition from hypothesis-tier to derived-result-tier must produce the missing derivation of the coordination penalty from a more fundamental AAT construct (the natural candidate is shared-intent mutual information between sub-agent models $M_t^{(A)}$ and $M_t^{(B)}$ under coupled-evidence regimes).
- **Dynamic-regime placement (added 2026-05-21 per Phase 5 cross-segment ripple).** Symbiogenic absorption is the **canonical structural mechanism for the R1 $\rightarrow$ R0 transition** on the dynamic-regime axis (per #disc-dynamic-regime-axis §”Transition asymmetry: descent is the default, ascent costs more”): pre-absorption, the composite operates in R1 equilibrium-regime under partially-opposing objectives between host and endosymbiont; post-absorption, the merged composite operates in R0 contraction-regime under the unified objective surface, with the (CM2) closed-form of #deriv-critical-mass-composition applying at the asymmetric-parameter limit. Symbiogenesis is therefore the self-driven *ascending* move on the dynamic-regime axis — paying the structural cost of objective-merger to exit the equilibrium-regime failure modes (saddle-Nash, multi-equilibria, last-iterate non-convergence) for the R0 contraction-regime guarantees. The saddle-node bifurcation analysis above (ρ_c threshold) is what makes the ascent *forced* rather than optional above critical environmental volatility.
- **Track E surface-back of catalog citations (2026-05-22) — Szathmáry 2015 major-transitions framework.** Adjacent literature surfaced 2026-05-22 from Track E catalog at ref/prior-art-analysis/08-composite-agency.md (Pillar 3): Szathmáry, E. (2015), *Toward major evolutionary transitions theory 2.0* (PNAS 112:10104)[†]. Substantial conceptual ancestry for the symbiogenic-absorption mechanism this segment formalizes: Szathmáry’s major-evolutionary-transitions (MET) framework names lower-level units becoming constrained by a higher-level unit, with progressively-reduced autonomy and “de-Darwinized” sub-units — exactly the asymmetric-parameter-limit structure (CM2)’s closed form formalizes mathematically. AAT does not claim to have discovered asymmetric absorption; the contribution is the *mathematical form* (the asymmetric-parameter limit reading of #deriv-critical-mass-composition + the dynamic-regime-axis R1 $\rightarrow$ R0 transition reading) that makes the symbiogenic process predictable rather than purely empirical. Other named conceptual ancestors (Maynard Smith & Szathmáry 1995 *The Major Transitions in Evolution*, Cambridge; Margulis 1981 *Symbiosis in Cell Evolution* — the canonical biological instance) are textbook/well-known and inherited from general scientific literacy rather than catalog-attributed.

Incidental audit gold (lift 2026-05-30). Cross-audit “wandering thoughts” / §14-ideation harvested from the de-novo auditors’ working dirs, deduplicated across substrates and lightly attributed. Orthogonal pedagogical / framing material, kept separate from the certified theory-fix findings (handled elsewhere — see report flags). **Coverage:** three dirs reached a digested reflection (526815, 829314, 849201) plus the §III composition-foundations batch (471203). The richest ideation here clusters around the weighted-Lyapunov $\mu \rightarrow 0$ sketch the segment’s Working Notes already carry — the gold is the *pedagogy around that math*, not new math. Substrate attribution inferred from voice where not explicit.

1. Candidate Brief prose / pre-prose.

- **The thermodynamic bargain — autonomy for persistence.** The single strongest framing here: symbiogenesis (and corporate acquisition, and mitochondrial endosymbiosis) is “a fundamental thermodynamic trade: **autonomy** ($\mu \rightarrow 0$) **in exchange for persistence** ($R \rightarrow R_h$).” The absorbed entity’s independent mismatch dynamics drop out of the joint stability equation in the limit $\mu \rightarrow 0$; what it *gains* is that it “no longer has to worry about the harsh environmental drift ρ or running out of runway (reserve R), because it is now sheltered inside the massive reserve R_h of the host.” This is offered as the intuition that explains why symbiogenesis is ubiquitous in both nature and capitalism (Gemini, AUDIT-WORKING-829314). A candidate Brief/Discussion anchor that makes the $\mu \rightarrow 0$ limit legible without symbols.

2. Candidate Discussion.

- **Symbiogenesis and peer coupling as parameter-limits of one equation.** The pedagogical reading of the existing $\mu \rightarrow 0$ Working Note: a “merger of equals” (peer coupling) weighs both parties’ stability dynamics roughly equally ($\mu \approx 1$) and must coordinate everything; an acquisition (symbiogenesis) sends the small party’s weight $\mu \rightarrow 0$ so its independent dynamics vanish from the joint stability argument. “Symbiogenesis isn’t a magical discontinuous leap, but a smooth parameter limit of the standard peer-coupling stability proof” (Claude, AUDIT-WORKING-849201; the merger-of-equals-vs-acquisition vignette at Gemini, AUDIT-WORKING-829314). Candidate Discussion prose making the segment’s sketch self-explaining. (*The auditors also urge promoting the $\mu \rightarrow 0$ limit from Working Notes into the Formal Expression — that is a strengthen/promotion question flagged in the report, not gold.*)
- **Why retained autonomy is not zero — the singular-perturbation reading.** “Why not just execute the startup’s founders and replace them with a script?” Because the host *wants* the fast, low-level local responses the endosymbiont was good at, without routing them through the slow macro-loop — so it “deliberately preserves a tiny sliver of the startup’s autonomy to handle high-frequency local disturbances that would violate the host’s slow tempo budget” (Gemini, AUDIT-WORKING-829314; same point, with a curiosity about deriving an *optimal* retained autonomy μ^* from local-disturbance spectral frequency ρ_{local}). Sharpens the segment’s existing (S-3)-“why not complete?” Working Note by tying it explicitly to temporal nesting / singular perturbation.

3. Follow-up items.

- **Optimal retained autonomy μ^* from local-disturbance spectrum.** A candidate quantitative extension of (S-3): derive the retained-autonomy weight μ^* as a function of the spectral frequency of the local disturbance ρ_{local} the endosymbiont absorbs (faster local disturbances $\rightarrow$ more retained autonomy). Connects to 4.9 (Gemini, AUDIT-WORKING-829314).
- **When is symbiogenesis favored over peer coupling? — transaction-cost framing.** Multiple substrates reached for transaction-cost economics (Coase / Williamson) as the analog: symbiogenesis is favored when the coordination overhead of maintaining separate objectives exceeds the integration cost (Claude, AUDIT-WORKING-849201; converges with the segment’s existing “Quantitative predictions” Working Note tying this to C_{coord}).

4. Readers often ask / wonder.

- **“What drives $O_e \rightarrow \mathcal{D}_e(O_h)$ — selection, bounded rationality, or design?”** A natural reader want for the most load-bearing open dynamic (S-1); the segment’s Working Notes already enumerate the three candidate drivers, and a fresh reader independently reached for the same set (Claude, AUDIT-WORKING-849201). Worth a one-line forward-pointer in the body so the reader knows the question is named.

5. Candidate figures.

- **Three simultaneous contractions, with a typed-ambiguity output.** Symbiogenesis drawn as three coupled contractions over time — (S-1) objective absorption, (S-2) function grafting, (S-3) autonomy reduction — with the output deliberately shown as a *typed ambiguity*: either a composite with a still-agentive subordinate sub-agent, or a host system with an integrated non-agentive component if autonomy falls below the agency-scope threshold (Codex/Claude, AUDIT-WORKING-526815, “diagram attempt”). The typing distinction it visualizes is also a real correctness concern flagged in the report (whether the absorbed entity is still an AAT sub-agent).

Belongs elsewhere.

- **Mitochondria / eukaryotic-cell as the canonical biological instance.** Repeatedly invoked (the eukaryotic cell as a parasite/predator that became a structurally necessary sub-component) (Gemini, AUDIT-WORKING-193847 under 12.2; Gemini, AUDIT-WORKING-829314). The segment already maps to mitochondria; this is convergence signal that the analogy lands, and the Margulis 1981 prior-art anchor is already in Working Notes.

“Citation surfaced 2026-05-22 from the Track E catalog at ref/prior-art-analysis/ (intermediate work artifacts that captured Pillar-style prior-art searches). Catalog has more verification support than raw Undermind synthesis but less than full primary-source reading. Verification queued with the BG2 cluster.

Chapter 13

Composition Machinery

[Gap] Discussion

13.1 *Formulation*: Composition Closure Criterion

The formal definition of when a group of interacting agents constitutes a *valid composite macro-agent*. The criterion: the closed-loop dynamics *approximately commute with coarse-graining* — projecting micro-states to macro-states and then running macro-dynamics yields approximately the same result as running micro-dynamics and then projecting. The **closure defect** ε^* is the minimum achievable mismatch between these two paths, taken over all admissible projections and macro-dynamics.

The construction accommodates a **timescale ratio** $K_c \geq 1$ — the number of micro-timesteps per macro-step — that lets the framework handle both lockstep composites ($K_c = 1$, synchronous-update) and composites that operate strictly slower than their sub-agents ($K_c \gg 1$, the regime 4.9 asserts as natural). When the ratio is large, the composite’s description need not — and should not — track the micro-trajectory between macro-boundaries; timescale abstraction is recovered.

The closure-defect framework applies *only* to sets that satisfy 12.2. Given scope-satisfaction, a set forms a *meaningful* composite agent when $\varepsilon^* \leq \varepsilon_{\max}$. Without scope-satisfaction, the infimum is still well-defined as a projection property, but the resulting “composite” has ill-defined purposeful substate and composition-level machinery (team persistence, composite tempo) does not apply.

The formulation specifies **admissibility constraints** (A1)-(A4) on the macro-dynamics — conditions the macro must satisfy to be AAT-shaped: (A1) the macro-state decomposes into model and purposeful components with a recursive update; (A2) a macro-mismatch signal is well-defined; (A3) macro-tempo is well-defined; (A4) the macro-correction function satisfies the sector condition (4.4). Together these make the composite a proper AAT agent rather than just a low-dimensional projection. Parallel constraints (P1)-(P3) govern the projection (information preservation, Lipschitz continuity, dimensionality reduction). A **bridge lemma** is the technical heart of the result: Part I’s sector-persistence template (G) applied at the composite level requires a *strictly stronger* contraction assumption on the macro-update map than (A4) alone supplies — the assumption is structurally motivated by (A4) but not formally derived from it. For different *Tiers* of sub-agent quality, Part I/II results transfer differently to the composite: Tier 1 transfers *exactly*; Tier 2 transfers locally with bounded degradation; Tier 3 requires per-domain verification.

The closure defect itself is the **effective disturbance** at the composite level — it enters the composite’s persistence inequality as the rate at which macro-tracking diverges from micro-truth, combining with external environment per-

turbations to set the composite's total disturbance load (13.2). This makes precise the *cost of imperfect composition*: a composite with high closure defect carries that defect as a perpetual disturbance source against its own persistence.

Formal Expression.

Let a system consist of N sub-agents interacting in a shared environment with state space $\mathcal{S}_{\text{env}}$. The micro-state, micro-observations, and micro-actions are:

$$X_{\text{micro},t} = \{(M_{i,t}, G_{i,t})\}_{i=1}^N \in \mathcal{X}_{\text{micro}} \quad (13.1)$$

$$o_{\text{micro},t} = \{o_{i,t}\}_{i=1}^N \in \mathcal{O}_{\text{micro}} \quad (13.2)$$

$$a_{\text{micro},t} = \{a_{i,t}\}_{i=1}^N \in \mathcal{A}_{\text{micro}} \quad (13.3)$$

The coupled micro-dynamics form an action-observation loop:

$$X_{\text{micro},t} \xrightarrow{\pi_{\text{micro}}} a_{\text{micro},t} \xrightarrow{E} (\Omega_{t+1}, o_{\text{micro},t+1}) \xrightarrow{f_{\text{micro}}} X_{\text{micro},t+1} \quad (13.4)$$

We constrain our search to an admissible class of projections $\Lambda \in \mathcal{P}_{\text{adm}}$ mapping micro to macro, an admissible class of macro-dynamics $(\pi_c, E_c, f_c) \in \mathcal{M}_{\text{adm}}$, and a **timescale ratio** $K_c \geq 1$: the number of micro-timesteps per macro-step. Micro-time is indexed by $t \in \{0, \dots, H\}$; macro-time by $m \in \{0, \dots, \lfloor H/K_c \rfloor\}$ with macro-step m corresponding to micro-timestep $t = mK_c$. The projection components then have type signatures:

- $\Lambda_x : \mathcal{X}_{\text{micro}} \rightarrow \mathcal{X}_c = (M_c, G_c)$ — pointwise, evaluated at macro-boundaries.
- $\Lambda_\Omega : \mathcal{S}_{\text{env}} \rightarrow \mathcal{S}_{\text{env},c}$ — pointwise, evaluated at macro-boundaries.
- $\Lambda_o : \mathcal{O}_{\text{micro}}^{K_c} \rightarrow \mathcal{O}_c$ — aggregates the window of K_c micro-observations between successive macro-boundaries.
- $\Lambda_a : \mathcal{A}_{\text{micro}}^{K_c} \rightarrow \mathcal{A}_c$ — aggregates the window of K_c micro-actions similarly (used on the observation-comparison side; the macro-policy π_c emits one macro-action per macro-step).

Aggregation is part of the projection specification — common choices are mean/sum (linear systems), first/last (event-rate projections), or task-specific sufficient statistics. When $K_c = 1$ the windows collapse, Λ_o, Λ_a reduce to pointwise maps, and micro- and macro-tempos coincide (synchronous-update composites). When $K_c \gg 1$ the composite operates on the quasi-steady-state output of its sub-agents, the regime that 4.9 asserts as the natural one for composition. K_c is part of the problem specification; it does not appear in (A1)-(A4) or (P1)-(P3) below, but it determines at what granularity closure is measured.

Let $\mathcal{D}_{\text{micro}}$ be the distribution of reachable trajectories generated entirely by the true micro-system over horizon H , and let $o_{\text{micro},(m-1)K_c:mK_c}$ denote the window of K_c micro-observations ($o_{\text{micro},(m-1)K_c+1}, \dots, o_{\text{micro},mK_c}$) between macro-boundaries $m-1$ and m (similarly for actions).

[*Definition (Composition Closure)*] We define the minimal achievable closure defect ε^* over the admissible classes as:

$$\varepsilon^* = \inf_{\Lambda \in \mathcal{P}_{\text{adm}}, (\pi_c, E_c, f_c) \in \mathcal{M}_{\text{adm}}} \left\| (\varepsilon_x, \varepsilon_a, \varepsilon_o) \right\| \quad (13.5)$$

where the expected component errors evaluated over true micro-trajectories $\tau \sim \mathcal{D}_{\text{micro}}$, measured **per macro-step**, are:

$$\begin{aligned}
\bullet \quad \varepsilon_x &= \mathbb{E}_\tau \left[\frac{1}{M_H} \sum_{m=1}^{M_H} \left\| \Lambda_x(X_{\text{micro}, mK_c}) - f_c(\Lambda_x(X_{\text{micro}, (m-1)K_c}), \Lambda_o(o_{\text{micro}, (m-1)K_c : mK_c})) \right\|_x \right] \\
\bullet \quad \varepsilon_a &= \mathbb{E}_\tau \left[\frac{1}{M_H} \sum_{m=1}^{M_H} \left\| \Lambda_a(a_{\text{micro}, (m-1)K_c : mK_c}) - \pi_c(\Lambda_x(X_{\text{micro}, (m-1)K_c})) \right\|_a \right] \\
\bullet \quad \varepsilon_o &= \mathbb{E}_\tau \left[\frac{1}{M_H} \sum_{m=1}^{M_H} \left\| \Lambda_o(E_{\text{obs}}(\Omega, a_{\text{micro}})|_{(m-1)K_c : mK_c}) - E_{c,\text{obs}}(\Lambda_\Omega(\Omega_{(m-1)K_c}), \pi_c(\Lambda_x(X_{\text{micro}, (m-1)K_c}))) \right\|_o \right]
\end{aligned}$$

where $M_H = \lfloor H/K_c \rfloor$ is the number of macro-steps in horizon H , and $E_{\text{obs}}(\Omega, a_{\text{micro}})|_{(m-1)K_c : mK_c}$ denotes the window of micro-observations the environment produces across the macro-step.

Units. $\varepsilon_x, \varepsilon_a, \varepsilon_o$ carry units of distance — per-macro-step norm errors. Multiplied by the macro-update rate ν_c (units $[\text{time}^{-1}]$), they yield the disturbance-rate units $[\text{distance}] \cdot [\text{time}^{-1}]$ demanded by the sector-persistence template's ρ_ξ (G) and by the dimensional accounting in 13.2. The previous per-micro-step formulation conflated ε^* 's granularity with ν_c 's — the macro-step formulation resolves the unit inconsistency.

$K_c = 1$ special case. When micro and macro update in lockstep, $M_H = H$, the observation/action windows collapse, and the formulas reduce to the synchronous form. This is the regime of tightly-coupled groups (distributed algorithms running in phase, ensemble filters where every member steps per tick) and is where the two-Kalman worked example below operates.

$K_c \gg 1$ case. When the composite lives at a strictly slower timescale than its sub-agents (the regime 4.9 asserts as natural), only micro-states at macro-boundaries enter ε_x , and macro-observations aggregate K_c micro-observations. The closure defect then measures how well the macro-dynamics predicts the *next* macro-state from the *previous* macro-state and the aggregated micro-observation window — not any intermediate micro-state. Timescale abstraction is recovered: the composite's description need not — and under $K_c \gg 1$ should not — track the micro-trajectory between macro-boundaries.

The closure-defect framework applies to sets that satisfy 12.2 — i.e., that form composites via at least one of the three alignment routes (shared objective, hierarchical derivation, mutual benefit). Given scope-satisfaction, a set forms a *meaningful* composite agent — distinguished from a multi-agent system whose low- ε^* projection is a happenstance of the environment rather than a reflection of composite coherence — when $\varepsilon^* \leq \varepsilon_{\text{max}}$. Without scope-satisfaction, the ε^* infimum is still well-defined as a projection property, but the resulting “composite” has ill-defined purposeful substate $G_c = (O_c, \Sigma_c)$ and composition-level machinery (team persistence, composite tempo) does not apply.

Admissibility constraints on macro-dynamics.

$\mathcal{M}_{\text{adm}}$ is the class of macro-dynamics (π_c, E_c, f_c) satisfying:

Formulation (macro-dynamics-admissibility).

(A1) AAT agent structure. The macro-state decomposes as $X_c = (M_c, G_c)$. The macro-update is recursive at macro-time m :

$$X_{c,m+1} = f_c(X_{c,m}, o_{c,m+1}) \quad (13.6)$$

The macro-policy is state-dependent: $a_{c,m} = \pi_c(X_{c,m})$. This ensures the macro-agent has the same structural components as a single AAT agent (1.1, 3.2, 3.3).

(A2) Macro-mismatch. A mismatch signal is well-defined:

$$\delta_{c,m} = o_{c,m} - \hat{o}_{c,m}(M_c, a_{c,m-1}) \quad (13.7)$$

where $\hat{o}_{c,m}$ is the macro-model's prediction of the next macro-observation. This ensures the macro-agent can detect prediction errors — the foundation of adaptation (3.4).

(A3) Macro-tempo. The macro-update has well-defined adaptive tempo:

$$\mathcal{T}_c = \sum_k v_c^{(k)} \cdot \eta_c^{(k)*} \quad (13.8)$$

where k indexes the macro-agent's observation channels, $v_c^{(k)}$ is the event rate, and $\eta_c^{(k)*}$ is the optimal gain (3.8, 3.6). This ensures the persistence condition is formulable at the macro level. The additive form carries 3.8's noise-independence scope at the *macro* channel level (A5): when macro channels aggregate cross-correlated micro-streams, the exact macro tempo is the joint form, and the additive expression deviates with sign — the composite-level consequences are worked in 13.2.

(A4) Bounded macro-correction. The macro-correction function satisfies the sector condition (4.4):

$$\delta_c^T F_c(\mathcal{T}_c, \delta_c) \geq \alpha_c \|\delta_c\|^2 \quad \text{for } \|\delta_c\| \leq R_c \quad (13.9)$$

with $\alpha_c > 0$ (positive correction rate) and $R_c > 0$ (finite reserve). This ensures the macro-agent's corrections work — they reduce mismatch rather than amplifying it, within the macro-model's capacity.

What (A1)-(A4) prevent. Without these constraints, the infimum over $\mathcal{M}_{\text{adm}}$ is trivially zero (any dynamical system can curve-fit micro-trajectories over a finite horizon). The constraints force the macro-dynamics to be a genuine AAT agent — with decomposed state, prediction errors, adaptive tempo, and stable correction — not an arbitrary function approximator. ε^* then measures the irreducible cost of representing multiple agents as one AAT agent.

What (A1)-(A4) do NOT require. Directed separation (6.3) is not part of the admissibility constraints. It is an additional structural property that some composites have and others don't (13.3). Results that depend on directed separation declare it as an additional assumption. Similarly, strategy structure ($G_c = (O_c, \Sigma_c)$ with a DAG) is not required — simpler goal representations are admissible.

Admissibility constraints on projections.

The admissible class of projections $\mathcal{P}_{\text{adm}}(\varepsilon_I, L)$ consists of projections $\Lambda = (\Lambda_x, \Lambda_o, \Lambda_a, \Lambda_\Omega)$ satisfying three conditions:

Definition (projection-admissibility, proposed from spike).

(P1) Information preservation. The projection retains at least a fraction $(1 - \varepsilon_I)$ of the predictive mutual information, evaluated across a macro-step:

$$I(\Lambda_x(X_{\text{micro}, (m-1)K_c}; \Lambda_o(o_{\text{micro}, (m-1)K_c : mK_c}) \mid \Lambda_a(a_{\text{micro}, (m-1)K_c : mK_c})) \geq (1 - \varepsilon_I) \cdot I(X_{\text{micro}, (m-1)K_c}; o_{\text{micro}, (m-1)K_c : mK_c} \mid a_{\text{micro}, (m-1)K_c : mK_c}) \quad (13.10)$$

The left side is the predictive information the macro-state has about the next macro-observation (aggregated over the K_c -window), conditioned on the macro-action. The right side is the micro-state's predictive information about the *same* macro-observation, conditioned on the same aggregated actions — both sides share a common target so the ratio is well-defined. The parameter $\varepsilon_I \in [0, 1]$ controls how much predictive power may be sacrificed by projection. This is the information bottleneck (Tishby et al. 1999) applied to the projection map. When $K_c = 1$, the windows collapse to singletons and (P1) reduces to the per-micro-step form previously stated.

(P2) Lipschitz continuity. Each component of Λ is Lipschitz-continuous with constant L :

$$\|\Lambda_x(X) - \Lambda_x(X')\|_{\mathcal{X}_c} \leq L \cdot \|X - X'\|_{\mathcal{X}_{\text{micro}}} \quad \forall X, X' \in \mathcal{X}_{\text{micro}} \quad (13.11)$$

(and analogously for Λ_o , Λ_a , Λ_Ω , each with its own Lipschitz constant). Without Lipschitz regularity, bounded closure defect does not imply bounded trajectory error — the bridge lemma requires this or something equivalent. When the projection is L -Lipschitz, the trajectory error bound becomes $L \cdot \varepsilon^* / \alpha_c$.

(P3) Dimensionality reduction. The macro-state space has strictly lower dimension than the micro-state space:

$$\dim(\mathcal{X}_c) < \dim(\mathcal{X}_{\text{micro}}) \quad (13.12)$$

This prevents the trivial identity projection, which achieves $\varepsilon^* = 0$ but is not a genuine abstraction.

Why three conditions. No single condition suffices alone. (P1) prevents degenerate projections (a constant map $\Lambda_x = c$ has zero mutual information). (P2) prevents pathological projections (discontinuous maps that amplify micro-errors into unbounded macro-errors). (P3) prevents trivial projections (the identity map satisfies (P1) and (P2) but achieves nothing). Together they constrain $\mathcal{P}_{\text{adm}}$ to projections that are informative, regular, and genuinely reductive.

The (ε_I, L) parameters are part of the problem specification, not derived quantities. The natural choice for many applications is $L = 1$ (non-expansive projection) and ε_I chosen as the minimum information loss compatible with genuine dimensionality reduction.

Independence from (A1)-(A4). The macro-dynamics admissibility (A1)-(A4) partially constrains the projection — (A1) requires Λ_x to preserve the M_c/G_c decomposition, and (A4) implicitly requires regularity through the sector condition. But (A1)-(A4) do not specify how much predictive information the projection must retain. A macro-agent with a very coarse projection can satisfy (A1)-(A4) while being a poor representation of the micro-system. The information-preservation condition (P1) fills this gap.

Bridge lemma: closure defect to trajectory error.

The bridge lemma is the sector-persistence template (G) applied with state variable $\xi = e_m = X_{c,m} - \Lambda_x(X_{\text{micro},mK_c})$ (trajectory error at macro-boundaries, between macro-evolved state and projected micro-state) and effective disturbance rate $\rho_\xi = \varepsilon^* \nu_c$ (closure defect per macro-step, multiplied by macro-update rate). The units are consistent by construction of the macro-step formulation above: ε^* is per-macro-step (distance) and ν_c is macro-step rate ($[\text{time}^{-1}]$), so ρ_ξ has the disturbance-rate units the template requires. The template yields directly:

Derived (bridge-lemma, from sector-persistence-template + A4 + contraction assumption).

$$\limsup_{m \rightarrow \infty} \|e_m\| \leq \frac{\varepsilon^* \nu_c}{\alpha_c} \quad (13.13)$$

This bound fits within the composite's sector-condition region iff $\varepsilon^* < \alpha_c R_c / \nu_c$, which is the persistence condition applied to the closure-error rate: if it fails, the macro-description diverges from micro-reality.

What the bridge lemma adds beyond the template. The template's precondition (T2) is the one-point sector bound. For the trajectory-error instantiation, the bound propagation requires a *strictly stronger* condition: the macro-update map $f_c(\cdot, o)$ must be contracting in its state argument (**incremental sector bound**, DA2'-inc — strong monotonicity of F_c), not merely one-point sector-bounded at each state. The per-step contraction factor is $\lambda = 1 - \alpha_c / \nu_c < 1$ (automatic when $\alpha_c < \nu_c$, true for any realistic agent that doesn't fully correct in a single step). This stronger condition is where the tier structure enters; see Epistemic Status for the Tier 1/2/3 taxonomy.

The sector condition (A4) does double duty — conditionally. (A4) ensures the macro-agent corrects external mismatch (single-agent persistence) AND — under the incremental sector bound — that the macro-description tracks

micro-reality (this bridge). Both are instances of the same template applied to different state variables with different effective disturbance rates.

What Is Derived vs. What Is Chosen.

Table 13.1

Property	Source	Strength
Closure-defect quantity ε^* as infimum over admissible classes	Definition built from (A1)-(A4), (P1)-(P3), and K_c	Formulation choice
Three-component decomposition into $\varepsilon_x, \varepsilon_a, \varepsilon_o$	Match to the three arrows of the action-observation loop (state, action, observation)	Formulation choice
Per-macro-step formulation; ε^* has units of distance-per-macro-step	2026-04-22 temporal coarse-graining repair (Finding A); forced by dimensional consistency with $\rho_\xi = \varepsilon^* \nu_c$ in G and 13.2	Derived (dimensional consistency)
Timescale ratio $K_c \geq 1$ as independent problem-specification parameter	Application-specific; $K_c = 1$ (synchronous) and $K_c \gg 1$ (4.9 regime) both admissible	Formulation choice
(A1) Macro AAT structure $X_c = (M_c, G_c)$ with recursive update	Import from 1.1 + 3.2 + 3.3	Formulation choice (requirement)
(A2) Well-defined macro-mismatch	Import from 3.4	Formulation choice (requirement)
(A3) Well-defined macro-tempo	Import from 3.8 + 3.6	Formulation choice (requirement)
(A4) Sector-bounded macro-correction	Import from 4.4	Formulation choice (requirement)
(A1)-(A4) as a requirement set render the infimum non-trivial (exclude curve-fitting macro-dynamics)	Consequence of restricting $\mathcal{M}_{\text{adm}}$ to genuine AAT agents	Derived (under the chosen requirement set)
(P1) Information-preservation constraint	IB Lagrangian-dual with source X_{micro} , relevance $O_{\text{micro}, t+1}$; formally connects to AA	Formulation choice (requirement)
(P2) Lipschitz continuity of Λ	Required for bridge-lemma trajectory-error bound (without it, bounded ε^* does not imply bounded trajectory error)	Formulation choice (requirement)
(P3) Strict dimensionality reduction	Rules out identity projection (which trivially achieves $\varepsilon^* = 0$ but is not abstraction)	Formulation choice (requirement)
(P1)-(P3) independent of (A1)-(A4)	(A1) and (A4) constrain Λ_x partially but do not fix information-preservation level (§"Independence from (A1)-(A4)")	Derived
Bridge lemma: $\limsup \ e_m\ \leq \varepsilon^* \nu_c / \alpha_c$	Sector-persistence template (G) applied with state $\xi = e_m$, disturbance rate $\rho_\xi = \varepsilon^* \nu_c$	Derived (conditional on incremental sector bound DA2'a-inc)
Persistence-as-closure-boundedness condition $\varepsilon^* < \alpha_c R_c / \nu_c$	Sector-condition region fits the asymptotic error ball	Derived (under A4 + DA2'a-inc)
Incremental sector bound (DA2'a-inc) is strictly stronger than (A4)'s one-point sector bound	Witness $F_d(\delta) = 2\delta + 3 \sin \delta$ (oscillatory, globally inward, locally non-monotone), Discussion §"DA2'a-inc is strictly stronger than (A4)"	Proved

Property	Source	Strength
Three-tier agent classification (Tier 1 / 2 / 3) for bridge-lemma applicability	Taxonomy induced by whether DA2'a-inc holds globally (T1), locally (T2), or must be verified per-domain (T3)	Formulation choice (classification)
Tier 1: bridge lemma derived for Bayesian updaters on exponential families, gradient descent on strongly convex losses, linear correction with positive-definite gain-observation product	Standard monotone-operator / Kalman analysis; $H^T H$ -norm contraction derivation in Discussion §"Bridge-lemma contraction"	Derived (conditional on DA2'a-inc + linear prediction, i.e., C2)
Tier 2: local contraction with factor degraded by $\kappa(D\delta)^2$	Extended-Kalman / locally-convex analysis	Derived (local)
Tier 3: contraction must be verified per-domain	No general result available for non-convex / discontinuous / non-mismatch-driven components	Discussion-grade
Composite (A4) from sub-agent properties: $\alpha_c \geq \min_i (\alpha_i - \Delta \mathcal{J}_i^{\text{cost}})$, $R_c \leq \min_i R_i$	Weakest-link bound (Discussion §"Deriving composite (A4) from sub-agent properties")	Derived (conditional on bounded coordination cost per 15.1)
Composite critical-mass inequality $(\alpha - C)R > \rho + \gamma \mathcal{J}$ (symmetric-matched-Tier-1) — sign-sensitive refinement of the weakest-link bound; recovers 15.1 (cooperative) and 15.2 (adversarial) as signed special cases	J	Derived (conditional on Tier 1 + matched-symmetric + Model D + coupling model C1/C2)
Composite persistence as (contraction) $\wedge$ (scope-satisfaction): $\kappa_c(U_O) > 0 \wedge$ 12.2	J (CM4); U_O enters multiplicatively on γ plus scope-gate, not additively	Derived (with (UO-mult) discussion-grade)
Asymmetric-limit connection to 12.3 (S-3) autonomy reduction	J §asymmetric-limit; weighted Lyapunov V_μ with $\mu \rightarrow 0$	Sketch (S-2 function transfer not yet formalized; promotable when 12.3 closes (S-2))
Meta-machine exact composition ($\varepsilon^* = 0$) for finite-state Moore machines	Miller 2022 §3.3; product automaton is the micro-dynamics by construction	Derived (external theorem)
Two-Kalman instantiation: $\varepsilon^* = 0$ at all ρ_{corr} for the means-only projection at steady state	Analytic steady-state calculation (Discussion §"Two-Kalman instantiation (closed form)")	Proved (closed form)
ε_x depends on both sub-agent unity and update-rule heterogeneity (ΔK)	Heterogeneous-gain Kalman case; 14.1	Derived (specific closed form for the two-Kalman case)
Mahalanobis norm as natural choice for estimation-type agents	Two-Kalman verification; Kalman gain minimizes expected squared Mahalanobis distance	Formulation choice (canonical for estimation; general-domain norm specification open)
N -agent scaling of ε^*	A "polynomial vs. exponential in N " framing is mis-typed: ε^* is a per-step residue, while "scales in N " is an accumulation-typed predicate	Resolved: $\varepsilon^*(N)$ is dimension-free-zero in the benign linear-Gaussian-stationary regime (all N , all coupling), graph-Laplacian-bounded with no exponential regime under compression (the consensus operator #deriv-critical-mass-composition already names), and order-incompatibility-invariant ($\leq S \log 2$, N -free) for strategy-DAG composition. No exponential regime exists. Full statement: the next two rows + Discussion
Computability of (P1) for nonlinear / non-Gaussian systems	Requires Monte Carlo or variational bounds; closed-form only for linear-Gaussian	Open

<i>Property</i>	<i>Source</i>	<i>Strength</i>
Principled setting of the ϵ_I threshold	No formal criterion currently; candidate tying to team size and coupling structure	Open
Strategy-DAG projection under Λ_x	Causal- τ -abstraction frame; 8.3 §"Composing heterogeneous strategy DAGs" (topology / SCC-condensation), 14.1 (fixed-topology credence instance)	Resolved (<i>conditional</i> ; one named sub-open Q1): the right object is a causal- τ -abstraction; exact Σ_c iff shared sub-orders jointly acyclic + AND/OR-compatible, else the minimal admissible Σ_c is the SCC-condensation with defect $\sum_{ C >1} \text{TC}(C) \leq S \log 2$ — bounded by shared-plan size, N-free. Beckers–Eberhardt–Halpern (τ, ϵ) -abstraction $\Rightarrow$ bounded ϵ_Σ (the <i>frame</i> , not term-for-term; the high-level distance remains AAT's to choose). Q1 (constant-tightness for non-deterministic intra-SCC coupling) is a named sub-open, not a reopening
Mori-Zwanzig zero-lag bound $\epsilon^* \geq \ Q_\Lambda UP_\Lambda\ _{\text{op}}$	Koopman/MZ projection (stationary coordinate-compatible case)	Derived (conditional on stationarity)
Full-kernel MZ bound on ϵ^*	Type mismatch (trajectory-accumulation vs. per-step)	Refuted at the ϵ^* -level; natural home is the bridge-lemma trajectory-error bound

The dividing line: (A1)-(A4) and (P1)-(P3) are *formulation choices of the requirement set* — alternative requirement sets are possible (e.g., entropy-preservation instead of MI-preservation; contraction bound in place of Lipschitz), and this one is chosen for parsimony, direct import from prior AAT segments, and downstream tractability. The *consequences* that follow under this requirement set — the non-triviality of the infimum, the bridge-lemma bound, the weakest-link composite-(A4) derivation, the two-Kalman $\epsilon^* = 0$ result — are Derived, with bridge-lemma-level results carrying the additional DA2'a-inc condition (which is proved strictly stronger than (A4) alone). The Tier 1/2/3 classification is itself a formulation, but the per-tier results are derived under tier-specific conditions. The genuinely-open questions are now the computability / principled-setting questions for (P1) and ϵ_I . The N -agent-scaling and strategy-DAG-projection rows are **resolved** (re-typed; the “poly-vs-exp” / “domain-specific-unresolved” framings were mis-typed — see the two rows above): there is no exponential-in- N regime, and N -scaling does not decide whether large teams are representable as single AAT agents.

Epistemic Status.

Conditional. Max attainable: conditional (formulation choice).

The closure defect ϵ^* is well-defined given (A1)-(A4), (P1)-(P3), the timescale ratio K_c , and specified norms. The macro-dynamics admissibility (A1)-(A4) and the projection admissibility (P1)-(P3) are independent formulation choices — (A1)-(A4) specify what “AAT-shaped macro-dynamics” means; (P1)-(P3) specify what “informative, regular, genuinely reductive projections” means. Both are needed for ϵ^* to be an infimum over a non-trivial, non-degenerate class. The timescale ratio $K_c \geq 1$ is a third, independent piece of problem specification: it sets the granularity at which closure is measured and the type signatures of Λ_o, Λ_a . K_c does not constrain the macro-dynamics class and is not derived — it is chosen from the timescale separation of the application (per 4.9). Both $K_c = 1$ (tightly-coupled composites) and $K_c \gg 1$ (hierarchically composed composites on their own timescale) are admissible. The projection admissibility conditions have a proposed definition with one exact instantiation (two-Kalman case at $K_c = 1$, worked in the Discussion); general-case computability of (P1) and the dependence of ϵ^* on K_c for a fixed application remain open.

The bridge lemma is *conditional* — not fully derived from (A4) alone. The precise additional condition is the **incremental sector bound** (DA2'a-inc): the correction function F_d must be *strongly monotone* ($\langle \delta - \delta', F_d(\delta) - F_d(\delta') \rangle \geq c_{\min} \|\delta - \delta'\|^2$ for all pairs), not just satisfy the one-point sector bound. This is strictly stronger than (A4) — a witness exists: the oscillatory correction $F_d(\delta) = 2\delta + 3 \sin \delta$ is globally inward-pointing but locally non-monotone (Discussion §"DA2'a-inc is strictly stronger than (A4)"). Three tiers of agents result:

- **Tier 1 (contraction proved):** Kalman filters, exponential-family Bayesian updaters, gradient descent on strongly convex losses, all linear corrections with positive definite gain-observation product. For these agents the bridge lemma is *derived* from (A4) + DA2'a-inc + linear prediction (C2), with contraction factor λ_{eff} from $\mathcal{O}$.
- **Tier 2 (local contraction):** Extended Kalman, locally convex gradients, nonlinear prediction models. Contraction holds locally with factor degraded by $\kappa(D\hat{o})^2$.
- **Tier 3 (independent verification):** Non-convex optimization, discontinuous/rule-based corrections, agents with non-mismatch-driven state components. Contraction must be verified per-domain.

The discrete-time recurrence argument itself is standard (Elaydi 2005); the contribution is identifying the incremental sector bound as the precise condition that bridges sector-bounded correction to full-update-map contraction. The full $H^\top H$ -norm contraction derivation is in the Discussion (§"Bridge-lemma contraction in the $H^\top H$ -norm").

The norm choices ($\|\cdot\|_{\mathcal{X}}$, $\|\cdot\|_{\mathcal{A}}$, $\|\cdot\|_{\mathcal{O}}$, and the combination norm) are load-bearing. For estimation-type agents, the Mahalanobis norm (weighted by inverse prediction-error covariance) is the natural choice — verified exactly in the two-Kalman case. For general domains, norm specification remains open.

Discussion.

This criterion replaces intuitive questions about “where the boundary of an agent is” with a functional test: does a macroscopic AAT description preserve the underlying micro-dynamics well enough to remain predictive and capable? The core requirement is an **approximate dynamical homomorphism** — the macro-dynamics approximately commute with the projection.

Relationship to 11. The Part I postulate requires that AAT’s machinery be scale-invariant — predictions at different levels of description must be compatible. This segment operationalizes “compatible” as “bounded closure defect under admissible coarse-graining.” The admissibility constraints ensure the macro-description is genuinely AAT-shaped, so the same persistence condition, the same tempo framework, and the same mismatch dynamics apply at the macro level with macro-level parameters.

Projection-behaviour facet of the stability certificate (5). Coarse-graining is a non-invertible projection of the certificate (**H**): the certificate-as-metric survives (the Schur complement of a positive-definite form is positive-definite) but the *dynamic* guarantee does not — the closure defect ϵ^* is the certificate’s projection-residue, equal to the Mori–Zwanzig zero-lag memory-commutator norm (derivation table above), zero exactly when the resolved subspace is invariant. Read through the spine, this is the question “does the agent’s measuring-stick survive being viewed coarsely”: its shape survives, its guarantee leaks by exactly the memory term. The composition-level identifiability floor (no common certificate across sub-agents; Liberzon, 5.1 Instance 3) is this same residue read as a boundary event — a *distinct* obstruction from the rank-collapse floors (a non-invertible projection, not a congruence), which is why the cross-sectional structure is several meta-patterns and not one.

The Jacobian-level observation. The incremental sector bound (DA2'-inc) used to prove the bridge lemma is mathematically equivalent to the contraction-metric condition (CT2) evaluated at the identity metric $M = I$, for continuously differentiable (C^1) correction functions F on convex domains. This equivalence (cf. Rockafellar & Wets 1998) means that AAT’s bridge lemma is a specialized application of generalized contraction theory. See #result-contraction-template for the full lifting of this condition to arbitrary Riemannian metrics.

Topology-indexed composition closures (via AI). `#result-contraction-template` lifts the bridge-lemma’s DA2’-inc condition into the contraction-metric framework of Lohmiller & Slotine 1998. Under its (CT1)–(CT3) preconditions, three topology-indexed closure results apply:

- **Parallel composition** preserves contraction under blockdiag metric at rate $\min(\lambda_1, \lambda_2)$ — recovers this segment’s weakest-link bound.
- **Cascade composition** ($\dot{x}_1 = f_1(x_1)$, $\dot{x}_2 = f_2(x_1, x_2)$) preserves contraction under lower-triangular weighted metric at rate bounded by $\min(\lambda_1, \lambda_2)$ up to coupling-gain adjustment (Slotine 2003 Thm 2). Handles hierarchical agent structures.
- **Negative feedback with bounded loop gain** (Slotine 2003 Thm 3) gives the heterogeneous critical-mass inequality $(\lambda_1 - C_1)(\lambda_2 - C_2) > k_{12}k_{21}/4$ via (CM2-M) in AI — extending J’s matched-symmetric (CM2) to heterogeneous sub-agents with different architectures and coupling strengths.

The Tier 1/2/3 bridge-lemma taxonomy of this segment maps cleanly to Tier 1M/2M/3M under AI (globally metric-contracting / locally metric-contracting / no globally valid metric). The observation that (CT2) at $M = I$ is *equivalent* to DA2’-inc means this segment has been carrying the Jacobian-level Euclidean contraction condition at the composite level all along; AI surfaces this at the single-agent level and adds the compositional theorems that generalize this segment’s closure content beyond the matched-symmetric case.

DA2’-inc $\equiv$ (CT2) at $M = I$ — structural-transparency equivalence. *[Derived, status: exact.]* For C^1 correction functions F on convex domains, the following three conditions are equivalent (standard monotone-operator theory; Rockafellar & Wets 1998 *Variational Analysis* Corollary 12.4; Nesterov 2004 *Introductory Lectures on Convex Optimization* §2.1.3 Theorem 2.1.9 — strong monotonicity $\Leftrightarrow$ symmetric-part of Jacobian uniformly positive-definite):

- (a) DA2’-inc (strong monotonicity): $(\delta - \delta')^\top (F(\delta) - F(\delta')) \geq c\|\delta - \delta'\|^2$ for all δ, δ' .
- (b) Jacobian-symmetric-part PD: $(\partial F / \partial \delta)_{\text{sym}} \succeq cI$ pointwise on the domain.
- (c) (CT2) at $M = I$ with $\lambda = c$: $-\partial F / \partial \delta - (\partial F / \partial \delta)^\top \leq -2\lambda I$.

Derivation. (a) $\rightarrow$ (b): take $\delta' \rightarrow \delta$ along direction v , divide by $\|\delta - \delta'\|^2$, take the limit. (b) $\rightarrow$ (a): integrate $v^\top (\partial F / \partial \delta)_{\text{sym}} v \geq c\|v\|^2$ along the segment from δ' to δ . (b) $\Leftrightarrow$ (c): direct algebraic identity.

Implication for AAT-internal derivability. Under this equivalence, the Euclidean sub-scope metric- α_1 cases of `#result-contraction-template` (Kalman-scalar, Euclidean strongly-convex-gradient, L2-regularized, linear-PD-symmetric) are AAT-internally derived from the DA2’-inc commitment already carried at the composite level — no new axiom required. What the (CT2)-at- $M = I$ framing adds is *visibility*: the Jacobian-level condition was implicit under the DA2’-inc name; surfacing the equivalence converts the Euclidean sub-scope from “theorem-imported from Lohmiller-Slotine” to “AAT-internally derived via a standard monotone-operator identity.” This is structural transparency, not new content. Non-Euclidean metric- α_2 cases (Fisher, Hessian, Lyapunov-metric) remain separately treated in `#result-contraction-template` — the equivalence is specific to $M = I$.

DA2’a-inc is strictly stronger than (A4): the witness. *[Derived, status: exact.]* The incremental sector bound (DA2’a-inc) does **not** follow from the one-point sector bound (A4) alone. A scalar witness exhibits a correction function satisfying the one-point bound while violating incremental monotonicity. Take

$$F_d(\delta) = 2\delta + 3 \sin \delta \quad \text{on } |\delta| \leq \pi. \quad (13.14)$$

One-point bound holds. Since $\delta \sin \delta \geq 0$ on $[-\pi, \pi]$ (the two factors share sign),

$$\delta F_d(\delta) = 2\delta^2 + 3\delta \sin \delta \geq 2\delta^2, \quad (13.15)$$

so the one-point sector bound (A4) holds with $\alpha = c_{\min} = 2$ on $\mathcal{B}_\pi$.

Incremental bound fails. In one dimension the incremental condition $(\delta - \delta')(F_d(\delta) - F_d(\delta')) \geq c_{\min}(\delta - \delta')^2$ is equivalent to $F_d'(\delta) \geq c_{\min}$ pointwise. Here $F_d'(\delta) = 2 + 3 \cos \delta$, and at $\delta = \pi$ this is $2 + 3 \cos \pi = 2 - 3 = -1 < 0$. So near $\delta = \pi$ there exist $\delta_1 < \delta_2$ with $F_d(\delta_1) > F_d(\delta_2)$, giving $(\delta_1 - \delta_2)(F_d(\delta_1) - F_d(\delta_2)) < 0$ — the incremental bound fails for *any* $c_{\min} > 0$.

The geometric reading: the one-point bound says “every correction points inward toward the equilibrium”; the incremental bound says “larger mismatches always receive larger corrections.” An oscillatory correction can be globally inward-pointing (one-point bound holds with margin) while *locally* the correction magnitude decreases as the mismatch grows (it wobbles) — which breaks incremental monotonicity. This is why the bridge lemma’s Tier-1 promotion carries the additional DA2’a-inc condition rather than resting on (A4): the witness shows the gap is real, not an artifact of proof technique.

Bridge-lemma contraction in the $H^\top H$ -norm (Tier 1). [Derived (bridge-lemma-contraction), conditional on (C1) incremental sector-Lipschitz, (C2) linear prediction, (C3) bounded gain.] The Tier-1 promotion of the bridge lemma — from “ f_c is assumed contracting” to “ f_c is contracting, derived” — rests on the following. Let the macro-update be mismatch-driven, $f_c(X, o) = X + \eta g(o - \hat{o}(X))$, with observation-space correction $F_d(\delta) = Hg(\delta)$, and assume:

- (C1) F_d is incrementally sector-Lipschitz: $\langle \delta - \delta', F_d(\delta) - F_d(\delta') \rangle \geq c_{\min} \|\delta - \delta'\|^2$ (DA2’a-inc, strong monotonicity) and $\|F_d(\delta) - F_d(\delta')\| \leq c_{\max} \|\delta - \delta'\|$;
- (C2) linear prediction $\hat{o}(X) = HX$, so $\delta - \delta' = H(X' - X)$;
- (C3) bounded gain $\eta < 2c_{\min}/c_{\max}^2$.

Then the full update map contracts in the $H^\top H$ -norm $\|v\|_{H^\top H}^2 = v^\top H^\top H v$:

$$\|f_c(X, o) - f_c(X', o)\|_{H^\top H} \leq \lambda_{\text{eff}} \|X - X'\|_{H^\top H}, \quad \lambda_{\text{eff}}^2 = 1 - 2\eta c_{\min} + \eta^2 c_{\max}^2 < 1. \quad (13.16)$$

Derivation. Applying H to the update difference $\Delta f = f_c(X, o) - f_c(X', o) = (X - X') + \eta[g(\delta) - g(\delta')]$ and writing $\Delta\delta = \delta - \delta' = H(X' - X)$,

$$H\Delta f = -\Delta\delta + \eta[F_d(\delta) - F_d(\delta')]. \quad (13.17)$$

Taking the squared Euclidean norm (which is the squared $H^\top H$ -norm of Δf):

$$\|\Delta f\|_{H^\top H}^2 = \|\Delta\delta\|^2 - 2\eta \langle \Delta\delta, F_d(\delta) - F_d(\delta') \rangle + \eta^2 \|F_d(\delta) - F_d(\delta')\|^2. \quad (13.18)$$

Applying (C1)’s incremental sector bound to the middle term and incremental Lipschitz bound to the last term:

$$\|\Delta f\|_{H^\top H}^2 \leq (1 - 2\eta c_{\min} + \eta^2 c_{\max}^2) \|\Delta\delta\|^2 = \lambda_{\text{eff}}^2 \|\Delta\delta\|^2, \quad (13.19)$$

and $\|\Delta\delta\|^2 = \|H(X' - X)\|^2 = \|X - X'\|_{H^\top H}^2$ closes the bound; (C3) makes $\lambda_{\text{eff}} < 1$. $\square$

This λ_{eff} is exactly the mismatch-dynamics contraction factor of $\bigcirc$ — the trajectory error e_m generates a virtual mismatch He_m contracted by the same mechanism that contracts actual mismatch, which is why the bridge bound $\limsup \|e_m\| \leq \varepsilon^* \nu_c / \alpha_c$ holds with no separate assumption once (C1)–(C3) do. When $H = I$ (direct observation) the norm is Euclidean and the factor is exactly λ_{eff} , recovering the scalar-Kalman case; for invertible H the Euclidean factor degrades by $\kappa(H)^2$. The classes for which (C1) holds — all linear corrections with positive-definite Hg , all exponential-family Bayesian updaters (strong monotonicity from log-partition convexity), all gradient descent on strongly convex losses

— are the Tier-1 agents; (C2)’s linear-prediction requirement is what separates Tier 1 (exact) from Tier 2 (local, factor degraded by the prediction Jacobian’s condition number squared).

The sector condition does double duty — conditionally. (A4) ensures the macro-agent’s corrections work (structural persistence), AND — given the additional contraction assumption (see Epistemic Status) — it ensures the macro-description tracks micro-reality (bridge lemma). Both uses address *structural persistence* in the sense of LEXICON.md — the capacity of the correction machinery, not the current operating point or the composite’s identity through time. The central insight is that structural persistence of the composite and faithfulness to its constituents are related through the same mechanism — contracting correction dynamics under bounded perturbation — but the bridge from sector-bounded correction to full-update-map contraction requires an assumption beyond (A4) that has been verified for estimation-type agents (Kalman) and remains open for general agents. When this assumption holds, the composite that persists in its own right also persists as a faithful representation of its constituents.

Deriving composite (A4) from sub-agent properties. If each sub-agent satisfies the sector condition with parameters (α_i, R_i) , and coordination costs are bounded by $\Delta\mathcal{F}_i^{\text{cost}}$ per agent (15.1), then the composite satisfies (A4) with:

$$\alpha_c \geq \min_i (\alpha_i - \Delta\mathcal{F}_i^{\text{cost}}) \quad (13.20)$$

$$R_c \leq \min_i R_i \quad (13.21)$$

This is a weakest-link bound (conservative but clean): each sub-agent’s effective correction rate is its own α_i degraded by its coordination overhead $\Delta\mathcal{F}_i^{\text{cost}}$, and the composite inherits the worst of these, while its reserve R_c cannot exceed the smallest sub-agent reserve. Cooperative coupling (15.1) can improve α_c beyond this bound by reducing effective disturbance. The key implication: (A4) is *verifiable* from micro-level properties — compute each sub-agent’s sector-condition parameters, estimate coordination costs, and check whether the composite has positive correction rate. No need to compute f_c directly.

Sign-sensitive refinement (derived). The sign-blind weakest-link bound is subsumed by J’s closed-form inequality $(\alpha - C)R > \rho + \gamma\mathcal{J}$ (CM2) in the matched-symmetric-Tier-1 case, where the right-hand side makes the effective disturbance explicit and the sign of the coupling γ (cooperative $\gamma < 0$ vs adversarial $\gamma > 0$) enters directly. (CM2) can yield composite persistence even when the weakest-link bound alone fails — when cooperative coupling reduces $\rho + \gamma\mathcal{J}$ below what the raw $\alpha - C$ margin would permit. 15.1 (cooperative) and 15.2 (adversarial) are recovered as signed special cases. The honest composite-persistence statement is the conjunction of contraction (CM3) and scope-satisfaction (12.2) — (CM4) — making explicit that composite persistence can fail in two qualitatively different ways: the composite fails to contract, or the composite was never a composite. The asymmetric-parameter limit $\alpha_2 \rightarrow 0$ via weighted Lyapunov V_μ formalizes 12.3’s (S-3) autonomy-reduction mechanism as a smooth deformation of (CM4), not a discontinuous regime change.

Connection to team-persistence. 15.1 derives persistence conditions for sub-agents in a cooperative-adversarial network. This segment provides the macro-level complement: the conditions under which the composite itself is a valid AAT agent. Together they close the loop: sub-agents persist individually (team-persistence) AND the composite is a meaningful macro-agent (composition closure with admissibility).

Constructive (A1)–(A4) via wrapping. The wrapping construction in #der-class-coercion-via-wrapping is a special-case instantiation where (A1)–(A4) hold *by construction* through the wrapper’s type signatures rather than by post-hoc verification of an admissible projection Λ . The wrapper builds explicit $X_W = (M_W, G_W)$ with belief-update and strategy-update maps that have specified type signatures; (A1) is automatic; (A2)–(A4) follow from wrapper-design constraints (commit to a prediction map; choose f_M from the Tier-1 belief-update class). The closure-defect ε^* in this special case has a different reading than the standard fidelity reading — the wrapper *deliberately* changes the underlying component’s behavior to enforce structural separation. Two distinct ε^* quantities appear in wrapper analyses: $\varepsilon_{\text{track}}^*$

(the standard fidelity quantity defined here, used in the bridge lemma) and $\epsilon_{\text{coerce}}^*$ (the wrapper-vs-component behavioral divergence, distinct from $\epsilon_{\text{track}}^*$ and not bounded by the bridge lemma). See #der-class-coercion-via-wrapping Discussion for the disambiguation.

Meta-machine: exact composition for finite automata (Miller 2022). When agents are finite-state automata (Moore machines), composition is *exact*: two machines with state sets S_1, S_2 interacting through their output/input channels form a **meta-machine** (Miller 2022, *Ex Machina*, §3.3) — itself a Moore machine with state set $S_1 \times S_2$, deterministic transitions (each machine’s output determines the other’s input), and a single starting state from the constituent starting states. The closure defect $\epsilon^* = 0$ trivially, because the product automaton *is* the micro-dynamics — there is no approximation. The meta-machine always falls into a cycle (finite states with deterministic transitions must eventually revisit a state), so two interacting automata produce an eventually-periodic joint behavior. The composition-closure framework becomes interesting when we ask for a *compressed* macro-description: can a smaller automaton (fewer than $|S_1| \times |S_2|$ states) approximate the meta-machine? This is where $\epsilon^* > 0$ and the admissibility constraints become non-trivial — (P3) requires genuine dimensionality reduction, and (P1) asks how much predictive information the compression preserves. Machine minimization (the theorem that every automaton has a unique minimized equivalent; Hopcroft et al. 2006) is a natural candidate for the optimal projection: the minimized meta-machine has the fewest states that reproduce the full joint behavior, achieving $\epsilon^* = 0$ with (P3) satisfied whenever the minimized machine is smaller than the product. See BD (planned) for the full instantiation.

Two-Kalman instantiation (closed form). The simplest nontrivial worked case: two scalar Kalman filters tracking a bivariate random walk with cross-correlated process noise $\text{Cov}(w_1, w_2) = \rho_{\text{corr}} q$ (common variance q , common observation noise r), each observing only its own component, no communication. This case sits at $K_c = 1$ — both sub-agents and the composite step together — so the macro-step formulation reduces to the pointwise form and the Λ_o aggregation is the identity. The natural projection keeps the state estimates and discards the covariance states (means-only projection, dimension 2 from micro-dimension 4). Each filter has the scalar steady state $P^* = \frac{1}{2}(-q + \sqrt{q^2 + 4qr})$, gain $K^* = P_{\text{pred}}^*/(P_{\text{pred}}^* + r)$, and marginal innovation variance $S^* = P^* + q + r$.

Although each filter is decoupled (agent i ’s update uses only o_i), the composite innovation vector $\delta_c = (\delta_1, \delta_2)$ is *cross-correlated* whenever $\rho_{\text{corr}} \neq 0$, because the environment couples the processes. The steady-state cross-covariance of the estimation errors $e_i = \omega_i - \hat{\omega}_i$ solves the fixed point $C_e = (1 - K^*)^2 C_e + \rho_{\text{corr}} q$, giving

$$C_e = \frac{\rho_{\text{corr}} q}{1 - (1 - K^*)^2} = \frac{\rho_{\text{corr}} q}{K^*(2 - K^*)}, \quad (13.22)$$

and the innovation cross-covariance is $\Gamma = \text{Cov}(\delta_1, \delta_2) = C_e + \rho_{\text{corr}} q$, so the composite innovation covariance is

$$\Sigma_\delta = \begin{pmatrix} S^* & \Gamma \\ \Gamma & S^* \end{pmatrix}, \quad \Gamma = \frac{\rho_{\text{corr}} q (1 + 2K^* - K^{*2})}{K^*(2 - K^*)}. \quad (13.23)$$

The per-step closure error of a macro-agent using gain K_c is $\epsilon_x^2 = \text{tr}((K^*I - K_c)\Sigma_\delta(K^*I - K_c)^T)$, minimized at $K_c = K^*I$ where it is exactly zero. So **at steady state the closure defect is exactly $\epsilon^* = 0$ for all ρ_{corr}** — the means-only projection perfectly represents the micro-dynamics because the Kalman gains have converged to constants and the discarded covariance state carries no information. All three projection conditions are verified exactly: $\epsilon_I = 0$ (the discarded covariance is a deterministic constant, so it carries zero conditional mutual information; (P1) holds with equality), $L = 1$ (the projection is a coordinate extraction with operator norm 1), and $\dim(\mathcal{X}_c) = 2 < 4 = \dim(\mathcal{X}_{\text{micro}})$. The admissibility conditions (A1)–(A4) hold for the macro-Kalman with $\alpha_c = K^*$, $R_c = \infty$, $\mathcal{T}_c = 2K^*$. The transient defect (before the gains converge) decays geometrically at rate $(1 - K^*)^2$ and is exponentially small after $O(1/|\log(1 - K^*)|)$ steps.

The “cost of independence” — the estimation performance lost by not exploiting cross-correlations — is a separate quantity, the **performance gap** $\Delta_{\text{perf}} = 2(P^* - p)$ between the independent composite and the optimal joint filter (diagonal

posterior covariance p). It is *not* a closure defect: it measures suboptimality relative to a centralized agent, not failure of the macro-description to track the micro-system. For small correlation it is second-order, $\Delta_{\text{perf}} \approx 2\rho_{\text{corr}}^2 q^2 r / (S^*)^2$ (the cross-gain is $O(\rho_{\text{corr}})$, and the information it extracts is $O(\rho_{\text{corr}}^2)$), growing monotonically to its maximum at $|\rho_{\text{corr}}| = 1$. The composition-closure framework diagnoses representability (“is this group a coherent AAT agent?” — yes), not optimality (“is this group as good as a centralized agent?” — no, when $\rho_{\text{corr}} \neq 0$). A group can be a perfectly valid composite agent while being suboptimal.

The first genuine $\varepsilon^* > 0$: purposeful sub-agents with Beta-Bernoulli strategy. The benign $\varepsilon^* = 0$ above turns on the discarded auxiliary state (the covariance) *converging* to a constant. Lift the sub-agents to purposeful agents, each carrying a single-edge strategy DAG with credence p_i updated Beta-Bernoulli on observed success/failure — $p_{i,t+1} = p_{i,t} + (\mathbb{1}[\cdot] - p_{i,t}) / (n_{i,t} + 2)$ — and the discarded auxiliary state is the observation count n_i , which *diverges* rather than converging. The edge-update gain $1/(n_i + 2) \rightarrow 0$ as $n_i \rightarrow \infty$. A macro-agent that discards n_i must use a fixed macro-gain η_c , and the per-step strategy closure error $\varepsilon_\Sigma = \mathbb{E}[|(1/(n_i + 2) - \eta_c)(\mathbb{1}[\cdot] - p_i)|]$ then grows without bound (fixed η_c cannot match a decaying gain) unless $\eta_c = 0$, at which point the macro-agent stops learning. So this is the first genuine source of $\varepsilon^* > 0$: a decaying-gain micro-dynamics is non-representable by fixed-gain macro-dynamics. It is structurally distinct from the Kalman case precisely in the convergence-vs-divergence of the discarded state — the Kalman covariance converges (defect vanishes at steady state); the Beta-Bernoulli count diverges (defect persists). The repair is to keep an effective-sample-size n_c in the macro-state (option (b)), restoring $\varepsilon^* = 0$ at the cost of a less aggressive projection — which signals that the purposeful substate carries more load-bearing dimensions than the epistemic substate. When $\varepsilon^* > 0$, the bridge lemma becomes load-bearing: the macro-strategy eventually goes stale and must be re-projected, the composition-level analog of 4.8.

Norm specification for estimation-type agents. The two-Kalman case identifies the Mahalanobis norm as the natural choice for agents whose primary function is state estimation. The state norm weights by the inverse prediction-error covariance: $\|X_c - X'_c\|^2 = (\hat{\omega}_c - \hat{\omega}'_c)^T (P_{\text{pred}}^*)^{-1} (\hat{\omega}_c - \hat{\omega}'_c)$. The observation norm weights by the inverse innovation covariance $S^{-1} = (P_{\text{pred}}^* + R)^{-1}$. The general principle: norms should weight by inverse uncertainty, so that differences in well-estimated components count more than differences in poorly-estimated ones. This is the norm the Kalman filter implicitly uses — the Kalman gain minimizes the expected squared Mahalanobis distance from truth. For domains without a natural covariance structure (discrete states, non-Gaussian models), the Euclidean norm remains the default.

Findings.

Composition-Closure Defect and Bridge Lemma. Brief: When does a group of interacting agents legitimately count as one larger agent? The closure-defect ε^* is the minimum residual prediction error from collapsing the micro-system into an AAT-shaped macro-description, taken over admissible coarse-grainings. The bridge lemma converts this prediction-level error into a trajectory-level guarantee: under sector-bounded macro-correction plus strong monotonicity of the macro-update, the macro-description tracks micro-reality with bounded asymptotic error $\varepsilon^* \nu_c / \alpha_c$. Composition is therefore not asserted by interface or by intuition — it is certified by an inequality, with explicit failure modes on both sides (the macro-dynamics fail to be AAT-shaped; the projection fails to be informative-and-regular; the macro-correction is not strongly monotone).

Impact: Closes the gap that Markov-blanket and active-inference accounts of nested agency leave open: those accounts give a statistical / thermodynamic boundary criterion but no control-theoretic loss bound on treating a collection as one agent. The closure-defect framework supplies the bound, attaches it to AAT’s existing sector-condition machinery, and routes Part I and Part II results across composite boundaries on Tier-1-classified sub-agents. Downstream this lets #deriv-critical-mass-composition derive the composite sector constant from sub-agent parameters in closed form, lets #der-tempo-composition give Brooks’s Law a formal expression in tempo units, and supplies the escape route (b) for #disc-identifiability-floor Instance 3. The criterion also draws an honest scope line: $\varepsilon^* = 0$ does not mean optimality (two non-communicating Kalman filters achieve $\varepsilon^* = 0$ while being suboptimal versus a joint filter); the framework diagnoses representability, not optimality.

Novelty Claim: *Claim differentiation* on bounded-loss composition as agent-boundary criterion. The bridge-lemma form is mathematically continuous with the approximate-information-state and influence-based-abstraction lines (Subramanian 2020; Congeduti 2020; Abel 2016; Taylor 2008), which prove approximation-loss results for compressed control under coarse-graining or aggregation. The differentiation is the *use* of such a bound as a criterion for when a multi-agent system can be re-described as a single agent — not as a control-policy compression result on a fixed agent — combined with explicit admissibility conditions (A1)-(A4) on the macro-dynamics that force the macro-description to itself be AAT-shaped, ruling out trivial curve-fits.

Related Work:

Table 13.2

<i>ASF concern</i>	<i>Prior-art language</i>	<i>Relationship / Positioning</i>
Statistical / thermodynamic boundary of a nested agent	Parr, da Costa & Friston 2019, “Markov blankets, information geometry and stochastic thermodynamics” <i>Phil. Trans. R. Soc. A</i> 378 (published 2019, found via Undermind 2026-04 Pillar 2)	<i>conceptual precursor</i> — supplies the boundary-and-nesting intuition; does not provide a coarse-graining error bound, a coordination-overhead theorem, or a strong-monotonicity condition for valid macro-agent treatment
Hierarchical-blanket nested agency with temporal depth	Kirchhoff, Parr, Palacios, Friston & Kiverstein 2018, “The Markov blankets of life” <i>J. R. Soc. Interface</i> 15 (published 2018, found 2026-04)	<i>conceptual precursor</i> — invokes synergetic ordering and self-organization at the compositional step but remains conceptual; no theorem bounding predictive loss of coarse-graining
Bound on control loss from compressed history	Subramanian, Sinha, Seraj & Mahajan 2020, “Approximate information state for approximate planning and reinforcement learning in partially observed systems” <i>arXiv:2010.08843</i> (published 2020, found 2026-04)	<i>closest mathematical neighbor</i> — proves a bound from predictive-compression error to control error in the same shape as the bridge lemma. ASF reuses the predictive-compression-to-control-loss mechanism and applies it at the coarser granularity of <i>agent boundary formation</i> , not single-agent state compression
Loss bound from approximate influence-based abstraction	Congeduti, Mey & Oliehoek 2020, “Loss Bounds for Approximate Influence-Based Abstraction” <i>arXiv:2011.01788</i> (published 2020, found 2026-04)	<i>closest mathematical neighbor</i> — closer than the Markov-blanket line to coordination-overhead mathematics. ASF differs in framing: Congeduti bounds policy-value loss under approximate cross-agent influence; the bridge lemma bounds <i>macro-state trajectory error</i> under coarse-graining
Near-optimality under state aggregation	Abel, Hershkowitz & Littman 2016, “Near Optimal Behavior via Approximate State Abstraction” <i>ICML</i> ; Taylor, Precup & Panangaden 2008, “Bounding Performance Loss in Approximate MDP Homomorphisms” <i>NeurIPS</i> (published 2016/2008, found 2026-04)	<i>adjacent literature</i> — single-agent MDP framing with aggregation bounds; same loss-from-coarse-graining mechanism, not framed as composite agency or organizational closure

<i>ASF concern</i>	<i>Prior-art language</i>	<i>Relationship / Positioning</i>
Exact reduction of decentralized control via common information	Nayyar, Mahajan & Teneketzis 2012, “Decentralized Stochastic Control with Partial History Sharing” <i>IEEE TAC</i> 58 (published 2012, found 2026-04)	<i>adjacent literature</i> — names the formal space where a lossy composition theorem would live but proves an exact reduction, not a closure-defect bound
Approximate dynamical homomorphism / model reduction	Mori-Zwanzig projection (1965); balanced truncation (Moore 1981); Koopman-operator analysis (Koopman 1931; modern review Mezić 2005)	<i>formal antecedent</i> — supplies the approximate-dynamical-homomorphism framing of the closure defect. The MZ zero-lag bound surfaces in this segment’s derivation table under stationarity. Adopted as standard machinery
Information bottleneck for the projection	Tishby, Pereira & Bialek 1999, “The information bottleneck method” <i>Allerton</i>	<i>formal antecedent</i> — admissibility condition (P1) is the IB Lagrangian-dual with source X_{micro} and relevance “next macro-observation given action”
Consistent abstractions of control systems	Pappas & Simic 2002, <i>Consistent abstractions of affine control systems</i> (IEEE TAC 47:745) ¹	<i>formal antecedent</i> — control-theoretic foundation for defining macro-systems via surjective coarse-graining. The macro-dynamics admissibility condition (A1) is operationalizing exactly Pappas-Simic’s surjectivity-of-coarse-graining requirement at the agent-state level
Compositional abstraction with explicit output-error bounds	Rungger & Zamani 2015, <i>Compositional Construction of Approximate Abstractions of Interconnected Control Systems</i> (IEEE TCNS 5:116) ²	<i>closest formal ancestor</i> — compositional abstraction with explicit output-error bounds is structurally the same shape as the closure-defect ε^* machinery this segment defines. The framework’s bridge-lemma quantitative-error treatment is a refinement of the Rungger-Zamani approach applied at the agent-boundary-formation layer rather than the interconnected-control-system layer
Markov-chain lumpability	Buchholz 1994, <i>Exact and ordinary lumpability in finite Markov chains</i> (J. Appl. Probab. 31:59) ³	<i>adjacent literature</i> — Markov-aggregation ancestor of the commute-with-coarse-graining criterion underlying admissibility (P1)–(P3). The lumpability framework distinguishes <i>exact</i> (commutativity-of-projection-and-transition) from <i>ordinary</i> (statistical-equivalence-of-aggregated-states) — the former is the strong-monotonicity condition the bridge lemma uses

Search Log:

- 2026-04 (*nominally comprehensive*): Two-search Undermind pass over composite agency, approximate information states, Markov-blanket nested agency, decentralized control, state abstraction, and influence-based abstraction. Verdict: *Conceptual Precursor* (High confidence). Markov-blanket line (Parr 2019; Kirchhoff 2018) positioned as conceptual precursor on nested agency without bridge-lemma mathematics; approximate-information-states / state-abstraction / influence-based-abstraction line (Sub20; Con20; Abe16; Tay08; Nay12) positioned as closest mathematical neighbors with shared predictive-loss-to-control-error mechanism but applied to single-agent or already-decomposed control problems rather than to agent boundary formation.

¹Citation surfaced 2026-05-22 from the Track E catalog at ref/prior-art-analysis/ (intermediate work artifacts that captured Pillar-style prior-art searches). Catalog has more verification support than raw Undermind synthesis but less than full primary-source reading. Verification queued with the BG2 cluster.

²Citation surfaced 2026-05-22 from the Track E catalog at ref/prior-art-analysis/ (intermediate work artifacts that captured Pillar-style prior-art searches). Catalog has more verification support than raw Undermind synthesis but less than full primary-source reading. Verification queued with the BG2 cluster.

³Citation surfaced 2026-05-22 from the Track E catalog at ref/prior-art-analysis/ (intermediate work artifacts that captured Pillar-style prior-art searches). Catalog has more verification support than raw Undermind synthesis but less than full primary-source reading. Verification queued with the BG2 cluster.

- 2026-04 (*intuition-only*, prior to comprehensive search): adjacent literatures expected to host prior art were dynamical-systems model reduction (Mori-Zwanzig, balanced truncation), active-inference Markov-blanket work, and PAC-MDP state-abstraction. The comprehensive search confirmed all three families and added the influence-based-abstraction and decentralized-control-via-common-information lines that intuition did not surface.

Working Notes.

- **Resolved:** $\mathcal{P}_{\text{adm}}$ now has a proposed definition. (P1)-(P3) above. Confirmed independent of (A1)-(A4) — the macro-dynamics admissibility partially constrains projections but does not specify information-preservation level. See `spikes/.integrated/spike-projection-admissibility.md` §5 for the analysis.
- **Resolved (2026-04-22): temporal coarse-graining gap.** The previous formulation summed $\varepsilon_x, \varepsilon_a, \varepsilon_o$ over micro-timesteps ($t = 1, \dots, H$), forcing synchronous micro-and-macro cadence. This contradicted 4.9 (composites naturally update on slower timescales, $\nu_{\text{level } n+1} \ll \nu_{\text{level } n}$) and conflicted with 13.2's dimensional accounting, which treats ε^* as a per-macro-step quantity and ν_c as the macro-update rate. Introducing the timescale ratio $K_c \geq 1$ and reformulating the defects per macro-step fixes both issues: Λ_o, Λ_a become window-aware (aggregating K_c micro-observations/actions), the sum runs over macro-steps m , and $\varepsilon^* \nu_c$ has consistent rate units by construction. $K_c = 1$ recovers the previous formula; $K_c \gg 1$ enables the timescale abstraction that was always part of the composition theory's intent. The bridge lemma and the sector-persistence-template instantiation are unchanged in substance — the template's state variable is now e_m at macro-boundaries rather than e_t at micro-timesteps, but the Lyapunov argument is identical. Option 3 of `audits/pending-findings-2026-04-21.md` Finding A; Option 2 (full Mori-Zwanzig equilibrium-residual form) remains a possible future refinement if one wants an explicit singular-perturbation framing.
- **Resolved: norm choices for estimation-type agents.** The Mahalanobis norm (inverse-covariance-weighted) is the natural choice for Kalman-type agents, verified exactly in the two-Kalman case. The general principle (weight by inverse uncertainty) extends to other estimation frameworks. For non-estimation agents (discrete states, non-Gaussian), norms remain domain-specific.
- **Open: computing (P1) for nonlinear/non-Gaussian systems.** The information-preservation condition requires conditional mutual information over the joint distribution of micro-states, observations, and actions. Tractable for linear-Gaussian systems (closed-form); requires Monte Carlo estimation or variational bounds for general systems.
- **Open: the right value of ε_I .** The information-preservation threshold is a free parameter. Too small ($\varepsilon_I \rightarrow 0$) excludes useful projections; too large ($\varepsilon_I \rightarrow 1$) admits degenerate ones. A natural candidate: ε_I comparable to the fractional information loss from adding one agent to the composite — tying it to team size and coupling structure. Formalizing this is open.
- **Resolved (2026-05-19): N -agent scaling of ε^* — the poly-vs-exp framing was mis-typed.** ε^* is a *per-step* residue; “scales polynomially or exponentially in N ” is an *accumulation*-typed predicate (the accumulation-type confound — see CHANGELOG 2026-05-19, `NOTATION.md` §“Accumulation typing”). Re-typed, it dissolves: dimension-free-zero (benign linear-Gaussian-stationary, all N , all coupling); graph-Laplacian-bounded with no exponential regime under compression (the consensus operator already in `#deriv-critical-mass-composition`); order-incompatibility-invariant $\leq |S| \log 2$ (N -free) for strategy-DAG composition. *History (not present truth): this note previously asserted an open “polynomial vs. exponential, coupling-structure-dependent” question and “the formal analog of the claim that very large teams cannot be treated as single agents”; that framing was the mis-type and is withdrawn — deleted and replaced when the re-typing landed, 2026-05-19, see CHANGELOG.* Reasoning trail: `spikes/spike-composition-scaling-N.md` (scoping) + the accumulation-type-confound trail (transient, $\rightarrow$.integrated/).
- **Partial: Mori-Zwanzig connection.** Under stationarity and coordinate-compatibility assumptions on the micro-dynamics, the Koopman-operator framing lifts Λ_x to a Hilbert-space projection P_Λ with $Q_\Lambda = I - P_\Lambda$. The MZ-optimal Markovian macro-dynamics is $f_c^{\text{MZ}} = P_\Lambda U P_\Lambda$ (the conditional expectation given the macro-state), providing a concrete benchmark against which admissibility constraints can be measured. When $f_c^{\text{MZ}} \notin \mathcal{M}_{\text{adm}}$, the per-step bound $\varepsilon^* \geq \|Q_\Lambda U P_\Lambda\|_{\text{op}}$ holds — a zero-lag memory-kernel bound. The bridge-lemma contraction on f_c in state space corresponds to spectral gap of $Q_\Lambda U$ in observable space (distinct quantities but related). *What does not close:* the full-kernel bound $\varepsilon^* \geq C \cdot \|K\|_{\ell^1}$ (sum-of-lags) is a type mismatch — ε^* is per-step; $\|K\|_{\ell^1}$ is trajectory-accumulation. Natural home for the full-kernel norm is the bridge lemma's trajectory-error bound, not ε^* directly. *Hard obstruction:* MZ's stationarity assumption fails for purposeful agents

with non-stationary auxiliary state (Beta-Bernoulli with diverging n) — exactly the cases where $\varepsilon^* > 0$ genuinely. Extending to these requires innovation-frame reformulation, not yet developed.

- **(P1) as IB Lagrangian-dual — resolved.** (P1) is the Lagrangian-dual of a standard IB constraint with source X_{micro} , compressed representation $\Lambda_X(X_{\text{micro}})$, relevance variable $o_{\text{micro},t+1} \mid a_{\text{micro},t}$, and $\beta(\varepsilon_I)$ the rate-distortion multiplier. See AA for the derivation and for the shared IB shape across AAT’s four compression operations (M_t, Σ_t , shared intent, Λ). (P2) Lipschitz continuity remains a separate analytic admissibility condition. (P3) dimensional reduction remains separate in the Gaussian case (the IB-optimal T at any finite β typically uses full support; categorical dimensional reduction is harder than any rate constraint). The “all admissibility is IB” slogan overclaims; the accurate slogan is “(P1) is IB; (P2) and (P3) compose with it as separate conditions.” Cross-instance unification is shape-sharing (U-medium), not a reduction to a single master problem (U-strong) — cross-instance theorems do not follow from the shared IB shape alone.
- **Two independent drivers of ε_x .** The non-degenerate Kalman case (heterogeneous gains $K_1^* \neq K_2^*$ with projection $\Lambda_X(\hat{\omega}_1, \hat{\omega}_2) = (\hat{\omega}_1 + \hat{\omega}_2)/\sqrt{2}$) shows ε_x depends on *both* sub-agent unity (process correlation, captured by U_M) and update-rule heterogeneity ($\Delta K = K_1^* - K_2^*$). $\Delta K = 0$ gives $\varepsilon_x = 0$ at every correlation; $\Delta K \neq 0$ gives $\varepsilon_x > 0$ even at perfect correlation. The closed form is $\varepsilon_x^2 = (\Delta K/2)^2 [S_- - C_{+-}^2/S_+]$ where $S_{\pm}$ are innovation variances in the $\pm$ rotation and C_{+-} their cross-covariance. Heterogeneity is not captured by the four unity dimensions in 14 as currently defined; this is a gap in that segment, not in this one — the closure defect correctly registers both drivers. See 14.1 for the full analysis.
- **Resolved (2026-05-19): strategy-DAG projection under Λ_X .** The right object is a causal- τ -abstraction of the joint micro-strategy (Beckers–Eberhardt–Halpern frame, adopted with citation; *not* term-for-term — BEH leaves the high-level distance application-dependent, so the norm-on-graphs choice remains AAT’s). Exact Σ_c iff shared sub-orders are jointly acyclic and AND/OR-compatible; otherwise the minimal admissible Σ_c is the SCC-condensation with closure defect $\sum_{|C|>1} \text{TC}(C) \leq |S| \log 2$ — bounded by shared-plan size, N -free. Status *conditional* with one named sub-open (Q1: constant-tightness for non-deterministic intra-SCC status coupling — a sub-open, not a reopening). *History (not present truth): this note previously listed the question as “deeply domain-specific and not resolved by (P1)-(P3)”;* resolved and replaced 2026-05-19 — see CHANGELOG. Derivation home: 8.3 § “Composing heterogeneous strategy DAGs” (the τ -abstraction typing, the SCC-condensation no-go, the Brooks’s-Law-floor corollary) and 14.1 (the fixed-topology credence-composition instance); reason-trail at spikes/.integrated/spike-strategy-dag-composition.md.
- **~~The bridge lemma’s contraction assumption — proving it from (A4) alone.~~ CHARACTERIZED 2026-04-06** (spikes/.integrated/spike-bridge-lemma-contraction.md). The contraction assumption cannot be proved from (A4) alone. The precise additional condition is the **incremental sector bound** (DA2’a-inc): the correction function F_d must be *strongly monotone* ($\langle \delta - \delta', F_d(\delta) - F_d(\delta') \rangle \geq c_{\min} \|\delta - \delta'\|^2$), not just satisfy the one-point sector bound. Strong monotonicity is strictly stronger than the one-point sector bound (counterexample: oscillatory corrections that are globally inward-pointing but locally non-monotone). Three agent tiers emerge: **Tier 1** (contraction proved — all Bayesian updaters on exponential families, all gradient agents on strongly convex losses, all linear corrections with positive definite gain-observation product); **Tier 2** (local contraction — nonlinear prediction models, with factor degraded by $\kappa(D\delta)^2$); **Tier 3** (independent verification required — non-convex optimization, discontinuous rules, agents with non-mismatch-driven state components). For Tier 1 agents, the bridge lemma is promoted from “conditional” to “derived (conditional on DA2’-inc + linear prediction).” The contraction factor equals λ_{eff} from O.
- The weakest-link structure ($\alpha_c = \min_i \alpha_i^{\text{eff}}$) is conservative. In practice, strong sub-agents may compensate for weak ones through cooperative coupling. A tighter bound would account for cross-agent compensation, likely through the cooperative disturbance reduction terms in 15.1.
- **A richer toy case** is needed: two purposeful agents (Part II) with strategy DAGs, cooperative communication, and a shared objective. This would exercise (A1) fully (including the G_c component) and test whether the admissibility constraints are tight enough to be useful without being so tight they exclude interesting composites.
- The approach is an approximate dynamical homomorphism condition, a standard tool in dynamical systems and model reduction (cf. Mori-Zwanzig projection, balanced truncation). The specific contribution is applying it to AAT’s closed-loop agent structure with sector-condition-based stability guarantees.

- **Categorical-cybernetics fit-check candidate (added 2026-05-22).** Capucci-Gavranović-Hedges-Rischel 2022 (*Towards Foundations of Categorical Cybernetics*, parametrised optics + selection functions) and Smithe 2024 (*Structured Active Inference*, generative models as “systems on an interface,” agents as controllers dual to generative models) supply a category-theoretic home for AAT’s composition machinery. The fit-check question for this segment specifically: does the closure-defect ε^* machinery admit a parametrised-optic formulation? Does ε^* have a natural reading as an optic morphism obstruction? Registered as `spikes/PROPOSED.md` Tier 2 first fit-check entry (2026-05-22), under the broader `spikes/PROPOSED-ADVANCED.md` Phase 3 #15 categorical-cybernetics 3-5-spike sequence. Long-horizon; high-leverage but substantial framework-foundation engagement. First-spike scope is *fit-check only* — if the fit is clean, subsequent spikes formalize per sub-question; if incompatible, recognition-tier citation acknowledgment in Discussion is the honest landing. Source-trail: `spikes/spike-enrichment-cluster2-2026-05-21/01-cluster-synthesis.md` Rank 4; landed via Track CR Phase 7.
- **Mori-Zwanzig memory-kernel relationship — upper-bound direction closes; named lower bound does not.** The closure defect ε^* has a tractable upper-bound relationship to a Mori-Zwanzig (MZ) memory kernel norm $\|K\|$ when AAT’s state-space projection Λ is composed with an MZ Hilbert-space projection P on a stationary observable measure π : under (i) stationarity of π under micro-dynamics, (ii) compatibility of Λ with P ’s range (the projected observables include the coarse-grained state coordinates AAT’s Λ retains), and (iii) bounded operator-norm of K , the upper bound $\varepsilon^* \leq C \cdot \|K\|$ holds for a constant C depending on α_c, ν_c . The *named target* lower bound $\varepsilon^* \geq C' \cdot \|K\|$ — which would make slow-decaying memory kernels an *obstruction* to small closure defect — does NOT close under the current AAT formulation without additional structural hypotheses. The obstruction is that Λ (state-space coarse-graining) and P (Hilbert-space projection on observables) are *different objects*, and a generic AAT coarse-graining does not give an orthogonal MZ-style projection on the observable Hilbert space; the lower bound would require either a stronger compatibility assumption between Λ and P or a separate orthogonality-of-residual argument that has not been worked out. The honest structural takeaway: MZ memory-kernel decay is a *sufficient* condition for small ε^* (via the upper bound direction) but is not *necessary*; AAT’s closure defect can be small for reasons that don’t reduce to MZ memory-kernel norms. Future work would either close the lower-bound direction under specific Hilbert-space hypotheses on the agent’s observable algebra, or settle the asymmetry as a structural feature distinguishing AAT’s coarse-graining from MZ projection-operator formalism.
- **Dynamic-regime placement (added 2026-05-21 per Phase 5 cross-segment ripple).** The (A1)–(A4) admissibility conditions and the closure-defect ε^* machinery this segment defines operate on **R0 contraction-regime** composites per `#dynamic-regime-axis`. R1 equilibrium-regime composites (under sub-scope α' of `#deriv-strategic-composition`) admit a parallel Lyapunov machinery via the A2'-analog transfer (potential or weighted-norm Lyapunov on joint deviation from Nash); R2 cyclic-distributional-regime composites (sub-scope β') admit only CCE set-convergence at $O(1/\sqrt{T})$ rate (Hart-Mas-Colell 2000) — neither inherits the bridge-lemma directly. The (A2) prediction map + (A4) sector condition presume a unique attracting fixed point in joint state space, which the R1 / R2 cases lack. Any segment citing `#form-composition-closure`’s machinery should also name the regime tier its sub-claim sits at.
- **Track E surface-back of catalog citations (2026-05-22).** Three Related Work entries added 2026-05-22 from Track E catalog at `ref/prior-art-analysis/08-composite-agency.md` (Pillars 1 + key anchor papers): Pappas-Simic 2002 (consistent abstractions of control systems — control-theoretic foundation for surjective coarse-graining); Rungger-Zamani 2015 (compositional abstraction with explicit output-error bounds — closest formal ancestor for ε^* machinery); Buchholz 1994 (Markov-chain lumpability — exact-vs-ordinary lumpability ancestor for the commute-with-coarse-graining criterion). All marked [`^cat-2026-05-22`] for verification-deferred attribution; primary-source verification queued in the BG2 cluster. Catalog citations had prior Pillar-style search support but were not verified by the current executor at landing time.

Incidental audit gold (lift 2026-05-30). Cross-audit “wandering thoughts” / §14-ideation harvested from the de-novo auditors’ working dirs, deduplicated across substrates and lightly attributed. Orthogonal pedagogical / framing / figure material, kept separate from the certified theory-fix findings (handled elsewhere — this segment drew a dense crop of correctness findings, the F108–F115 cluster, all routed to the findings track; see report flags). **Coverage:** four dirs reached a digested reflection (193847, 526815, 829314, 849201) plus the §III composition-foundations batch (471203). Multiple substrates called this “the strongest Part III segment” / “the theoretical peak of the framework” — the highest-engagement composition segment. Substrate attribution inferred from voice where not explicit.

1. Candidate Brief prose / pre-prose.

- **“When is a group of things a single thing?”** The closure defect ε^* stated as the mathematical answer to the ancient reductionism-vs-emergence question — “a society *is* a valid macroscopic entity if and only if $\varepsilon^* < \alpha_c R_c / \nu_c$.” If the inequality holds, the macro-description (sociology, macro-economics) tracks reality well enough to persist and “you don’t need to track every individual”; if it fails, “the society is just a chaotic multi-agent system” and macro-modeling of *that* group is an illusion (Gemini, AUDIT-WORKING-829314; the same “exact mathematical test for whether a group can be modeled as a single agent” framing at Claude, AUDIT-WORKING-849201 and Claude, AUDIT-WORKING-773921). The single strongest Brief/Discussion-opener candidate in this chapter.
- **Representability, not optimality.** The two-Kalman result ($\varepsilon^* = 0$ for two non-communicating filters tracking correlated walks, at *all* correlations) glossed as: “you can perfectly mathematically model a stupid organization as a single stupid macro-agent” — ε^* measures whether the group *acts like a single thing*, not whether it *acts well*; the “cost of independence” is a performance gap, not a closure defect (Gemini, AUDIT-WORKING-193847; Gemini, AUDIT-WORKING-829314; sharpest at Claude, AUDIT-WORKING-471203 — “the composition-closure framework diagnoses representability, not optimality ... exactly the right scope-honesty move”). Already a Working-Notes / Discussion theme; this is converged plain-language phrasing for it.

2. Candidate Discussion.

- $K_c \gg 1$ is **“the math of management.”** The timescale ratio glossed pedagogically: at $K_c = 1$ the manager updates at worker-speed and must observe every action in real time — micromanagement, which violates 4.9; at $K_c \gg 1$ the manager looks only at *aggregate* output over the window, which is Auftragstaktik (mission command), and the math proves this is exactly what enables genuine dimensionality reduction (P3) and information compression (P1) — “the manager doesn’t need to know *how* the workers achieved the goal, only *that* they did” (Gemini, AUDIT-WORKING-193847). A candidate Discussion vignette that makes the K_c / window-aware-projection machinery legible. The corporate-quarterly-update version (“the closure defect measures whether the corporation’s Q2 state can be predicted from its Q1 state, ignoring all the daily micro-fluctuations”) at Gemini, AUDIT-WORKING-829314 is the same point.
- **The Bridge Lemma converts prediction error into a survival bound.** Framing for the lemma’s payoff: it “converts abstract prediction error (defect) into physical trajectory bounds (survival)” — the mechanism by which “macro-economics or sociology can exist as predictive sciences distinct from micro-physics/psychology, provided the coarse-graining Λ is admissible” (Gemini, AUDIT-WORKING-193847; Gemini, AUDIT-WORKING-829314). Candidate Discussion gloss on *why* the bridge lemma is the chapter’s crown jewel.
- **The (A4)-insufficient / DA2’-inc-required point, plainly.** Several substrates independently praised the explicit proof that the one-point sector bound (A4) is *insufficient* and the strictly-stronger incremental sector bound (strong monotonicity, contraction between *any two trajectories* not just toward the origin) is required — with the Tier 1/2/3 classification (Kalman / convex = Tier 1; non-convex NN = Tier 3) read as “flawless dynamical-systems theory” and the DA2’-inc $\equiv$ Lohmiller-Slotine (CT2)-at- $M = I$ equivalence read as converting an *imported* theorem into an *AAT-internally-derived* consequence — “structural transparency, not new content” (Claude, AUDIT-WORKING-849201; Claude, AUDIT-WORKING-773921; Claude, AUDIT-WORKING-471203). Worth preserving the trajectory-vs-origin contraction distinction as the plain-language hook for why (A4) alone doesn’t suffice.

3. Follow-up items.

- **Most consequential modern composites are Tier 3.** A scope-honesty observation worth surfacing: transformer attention is non-convex, prediction models nonlinear, update dynamics not naturally mismatch-driven — so most LLM-based composite agents land in Tier 3, where the bridge lemma’s *global* guarantee does not apply (Tier 2/3 results apply with caveats) (Claude, AUDIT-WORKING-471203). Downstream theorems relying on the bridge lemma should constrain scope to Tier 1 or acknowledge local/heuristic status for Tier 2/3 (Claude, AUDIT-WORKING-849201 raised this as a “critical finding potential” to police downstream).
- **(P1) computability is narrow.** The information-preservation condition is tractable in closed form only for linear-Gaussian systems; general systems need Monte Carlo or variational bounds, and the two-Kalman case is the only fully verified instance

(Claude, AUDIT-WORKING-471203). Already an open Working Note; converged-on as a practical-applicability limit worth stating where a reader would reach for it.

4. Readers often ask / wonder.

- **“How does ε^* scale with N ? Can a hierarchy push $N \rightarrow \infty$ with bounded defect?”** A natural reader question raised by three substrates (Gemini, AUDIT-WORKING-193847 — wondering polynomial vs exponential; Claude, AUDIT-WORKING-849201; Claude, AUDIT-WORKING-471203 — “the structural test of whether very large teams can be treated as single agents at all”). The segment’s Working Notes now *answer* this (the poly-vs-exp framing was an accumulation-type confound; re-typed it dissolves to dimension-free-zero / graph-Laplacian-bounded / $\leq |S| \log 2$). Strong signal that the resolved answer should be surfaced in the body where the question naturally arises — fresh readers keep asking it.

5. Candidate figures.

- **The approximate-commuting square.** The closure criterion drawn as a square that fails to close: true micro-dynamics across the top, projection Λ down the side, macro-dynamics across the bottom — and ε^* is the gap by which the square fails to close. Side gates for projection admissibility (P1–P3) and macro-agent admissibility (A1–A4); a bridge arrow from per-step defect to bounded trajectory error under incremental contraction (Codex/Claude, AUDIT-WORKING-526815, “diagram attempt”). The commuting-square is the standard picture for an approximate dynamical homomorphism and would make “composition is a bounded-loss coarse-graining, not a vibe or an interface boundary” immediate.

Belongs elsewhere.

- **Consciousness-architecture — abstraction requires a slower self-clock.** The $K_c \gg 1$ reading extended to interiority: a mind’s macro-state (“its high-level ‘I’”) must run at a slower clock than its perceptual/motor sub-agents — if it tries to track every token or motor twitch at $K_c = 1$, the projection cost overwhelms the system even when ε^* is low. “True consciousness requires abstraction, and abstraction requires ignoring the high-frequency jitter of your own sub-components” (Gemini, AUDIT-WORKING-193847). Aspirational reach pointing at 03-llm-core/ / 04-eli-core/ (and at 4.9), not at this formulation segment.
- **Categorical-cybernetics home for ε^* .** The closure defect as a candidate “optic morphism obstruction” in the Capucci-Gavranović-Hedges-Rischel 2022 parametrised-optics framework — “if AAT can be fully expressed in Category Theory, it becomes substrate-independent at the deepest possible mathematical level” (Claude, AUDIT-WORKING-773921). Already registered as a spikes/PROPOSED.md Tier-2 fit-check in this segment’s Working Notes; convergence signal that a fresh reader independently reached for the same home.

13.2 *Derived (conditional): Derived: Composite Tempo Inequality*

The adaptive tempo of a composite agent is *bounded above* by the **joint tempo** of its sub-agents’ pooled observation streams — not, in general, by the sum of their individual tempos. The gap between the joint potential and the realized composite tempo is the **coordination overhead** C_{coord} — tempo consumed by *internal reconciliation* rather than external mismatch correction. The relation between the joint baseline and the additive accounting $\sum_i \mathcal{F}_i$ is *signed*, inheriting AS’s two-sided deviation across agents: equal under cross-agent noise independence, smaller under common-source noise (the echo-chamber team), and **larger under anti-correlated noise — a composite can genuinely outrun the sum of its parts** (triangulating team), realizable at exactly zero closure defect by fusion as simple as averaging.

The intuition, corrected. Tempo is a rate of mismatch correction, and the composite’s corrective capacity is set by the *information in its parts’ pooled streams* — which is not the sum of the parts’ individual capacities, because cross-agent noise structure is capacity no individual part has (anti-correlated errors cancel under fusion) or capacity the additive accounting double-counts (a shared upstream source). When closure defect is positive, sub-agents additionally spend tempo correcting *internal* mismatches generated by other sub-agents (Alice changes an API; Bob must update his model) or taking actions that counteract each other. That fraction is the coordination overhead and is *unavailable for external mismatch correction*.

The derivation connects closure defect to coordination overhead through *dimensional accounting*. The closure defect ε^* (units: distance) times the macro-update rate ν_c (units: inverse time) yields an internal disturbance rate matching the units of environmental disturbance. This internal mismatch must be corrected by the composite's correction capacity, and the correction consumes tempo that would otherwise address external mismatch. Conversion from disturbance rate to tempo-equivalent requires dividing by the task-adequacy distance scale $\|\delta_{\text{critical}}\|$ — mirroring the persistence-condition form used at the single-agent level. The composite therefore faces total *effective disturbance* combining external and closure-defect contributions, with coordination overhead bounded below by $C_{\text{coord}} \geq \varepsilon^* \nu_c / \|\delta_{\text{critical}}\|$.

This makes Brooks's Law precise. Adding sub-agents increases potential aggregate tempo $\sum \mathcal{T}_i$, but if it also increases closure defect (more coordination links to maintain), the realized composite tempo can *fall* — adding people to a late project makes it later when the closure-defect-driven tempo penalty from the additional links exceeds the new member's tempo contribution. The result is labeled *sketch* rather than exact because the closure-defect-to-coordination connection depends on the bridge lemma in 13.1 (which requires a contraction assumption on the macro-update map that is structurally motivated by (A4) but not formally derived from it). The remaining gap: whether the lower bound is tight or whether additional coordination costs (negotiation, synchronization, conflict resolution) contribute beyond the closure-defect mechanism.

Formal Expression.

This segment instantiates the sector-persistence template (G) at the composite level with state variable $\xi = \delta_c$ (composite mismatch) and a decomposed effective disturbance $\rho_{\text{eff}} = \rho_{\text{ext}} + \varepsilon^* \nu_c$ — external environment plus closure-defect contribution from 13.1's bridge lemma. The template supplies the Lyapunov machinery; this segment's distinctive content is the tempo-equivalent form of the coordination overhead lower bound $C_{\text{coord}} \geq \varepsilon^* \nu_c$.

Let $\mathcal{T}_i$ be the adaptive tempo of sub-agent i within a composite group of N agents. Let $\mathcal{T}_c$ be the adaptive tempo of the composite macro-agent A_c defined by an admissible coarse-graining Λ .

Let $\mathcal{T}_{\text{joint}}$ denote the joint tempo of the sub-agents' pooled observation channels — the exact object of AS ($U_c \cdot \mathbf{1}^T \Sigma_n^{-1} \mathbf{1}$ in the Fisher-local regime, with Σ_n the pooled cross-agent noise covariance).

[Derived (Joint-Bounded Tempo, Conditional on the Fisher-local regime)] For any admissible composite agent A_c , the composite tempo is bounded by the joint tempo of its parts' pooled streams: $\mathcal{T}_c \leq \mathcal{T}_{\text{joint}}$. No admissible macro-dynamics extracts more corrective information per unit time than the pooled micro-observations carry (data processing at the Fisher level).

[Definition (Coordination Overhead)] The **coordination overhead penalty** C_{coord} is the difference between joint potential and realized macro-tempo: $C_{\text{coord}} := \mathcal{T}_{\text{joint}} - \mathcal{T}_c$

[Derived (Signed additive-baseline deviation, from AS applied across agents)] The additive accounting relates to the joint baseline with indefinite sign: $\mathcal{T}_{\text{joint}} \lesseqgtr \sum_{i=1}^N \mathcal{T}_i$, with equality iff the cross-agent noises are independent (plus the measure-zero cancellation set of AS), $<$ under common-source cross-agent noise, and $>$ under noise-cancelling (anti-correlated) structure. The historical form of this segment's claim, $\mathcal{T}_c \leq \sum_i \mathcal{T}_i$, is therefore **conditional on cross-agent noise independence** and false in general — with a closure-exact witness:

[Derived (Super-additive witness, exact)] Two sub-agents observing $o_i = \theta + n_i$ with $\text{Var}(n_i) = r$ and $\text{Corr}(n_1, n_2) = \rho < 0$, each running its own optimal filter; composite defined by the averaging projection $\Lambda_x = \frac{1}{2}(m_1 + m_2)$, macro-observation $\bar{o} = \frac{1}{2}(o_1 + o_2)$, and the recursive running-mean macro-update. The macro-update commutes with the micro-dynamics *identically* ($\varepsilon^* = 0$ exactly; (A1)–(A4) and (P1)–(P3) verified), and $\frac{\mathcal{T}_c}{\sum_i \mathcal{T}_i} = \frac{1}{1+\rho} > 1$, unboundedly as $\rho \rightarrow -1$ — tempos evaluated in the matched-uncertainty (Fisher-inflow) convention of AS; the convention is load-bearing (Epistemic Status). A faithful composite of a triangulating team outruns the sum of its parts; symmetrically, at $\rho > 0$ the same witness has $\mathcal{T}_c < \sum_i \mathcal{T}_i$ at zero closure defect and zero coordination activity — the additive baseline,

not the composite, is what fails. Full derivation + numerics: spikes/spike-composite-tempo-synergy-2026-07-16.md (Working Notes).

Epistemic Status.

Conditional. Max attainable: exact. Component strengths: the joint bound $\mathcal{T}_c \leq \mathcal{T}_{\text{joint}}$ is Fisher-level data processing, exact within the first-order tempo–Fisher regime of AS ((R1)–(R3) of AR — the regime where tempo is a rate); the signed additive-baseline deviation is AS’s theorem applied across agents (same tier); the super-additive witness is *exact* — commutation is an identity, and the tempo ratio $1/(1 + \rho)$ is closed-form, numerically confirmed (spike, Working Notes). The ratio is stated in the matched-uncertainty (Fisher-inflow) convention — comparing information inflow at equal uncertainty, as throughout AS; the convention is load-bearing because static estimation makes literal own-trajectory fractional-contraction tempo estimator-blind. Independent verification probed convention-robustness: under dynamic θ (random-walk environment, each filter at its own steady state — the own-uncertainty reading) the witness’s ratio becomes $\sqrt{1/(2(1 + \rho))}$ with super-additivity threshold $\rho < -1/2$ rather than $\rho < 0$ — the refutation of the additive bound survives every tested reading; only the exact ratio and threshold are convention-tied. The historical sub-additive claim $\mathcal{T}_c \leq \sum_i \mathcal{T}_i$ survives only on the cross-agent noise-independence scope; its former grounding (“composition cannot create corrective capacity out of nothing”) was refuted — cross-agent noise structure *is* corrective capacity that no individual part has, and the additive accounting cannot see it. The $\varepsilon^* \rightarrow C_{\text{coord}}$ connection remains *sketch-grade*: it depends on the bridge lemma in 13.1 (contraction assumption DA2’a-inc, structurally motivated by (A4) but not derived from it), and whether the lower bound is tight — or additional process costs (negotiation, synchronization, conflict resolution) contribute beyond the closure-defect mechanism — is open.

Intuition. Tempo is a rate of mismatch correction. When $\varepsilon^* = 0$ (exact closure) and the composite’s fusion is information-lossless, $\mathcal{T}_c = \mathcal{T}_{\text{joint}}$ — and $\mathcal{T}_{\text{joint}}$ coincides with $\sum \mathcal{T}_i$ exactly when cross-agent noises are independent. When $\varepsilon^* > 0$, sub-agents spend part of their tempo correcting *internal* mismatches generated by other sub-agents (Alice changes an API; Bob must update his M_t) or taking actions that counteract each other. That fraction is C_{coord} and is unavailable for external mismatch correction.

The closure-defect to coordination-overhead connection.

The closure defect ε^* introduces internal mismatch at rate $\varepsilon^* \nu_c$ — units of [distance] · [time^{−1}], a disturbance rate matching the units of environmental ρ . This internal mismatch must be corrected by the macro-agent’s correction capacity, and the correction consumes tempo that would otherwise address external mismatch.

Derived (coordination-overhead-bound, from composition-closure bridge lemma).

Dimensional accounting. Three quantities appear in the composite persistence analysis, each with its own units:

Table 13.3

Quantity	Units	Role
Tempo $\mathcal{T}, \mathcal{T}_c, \mathcal{T}_c^{\text{ext}}, C_{\text{coord}}$	[time ^{−1}]	Rate of mismatch correction
Disturbance rate ρ, ρ_{eff}	[distance] · [time ^{−1}]	Rate of mismatch injection
Closure defect ε^*	[distance]	Per-step norm error
Critical scale $\ \delta_{\text{critical}}\ $	[distance]	Task-adequacy boundary

The conversion between a disturbance rate and a tempo-equivalent requires division by a distance scale, mirroring the persistence condition form $\mathcal{T} > \rho / \|\delta_{\text{critical}}\|$.

Effective disturbance. The macro-agent faces total effective disturbance combining external and closure-defect contributions:

$$\rho_{\text{eff}} = \rho_{\text{ext}} + \varepsilon^* \nu_c \quad (\text{units: } [\text{distance}] \cdot [\text{time}^{-1}]) \quad (13.24)$$

Coordination overhead (tempo-equivalent). The fraction of macro-tempo consumed by closure correction, expressed as a tempo:

Derived (coordination-overhead-tempo-form).

$$C_{\text{coord}} \geq \frac{\varepsilon^* \nu_c}{\|\delta_{\text{critical}}\|} \quad (\text{units: } [\text{time}^{-1}]) \quad (13.25)$$

The closure error rate is converted to a tempo penalty by normalizing against the task-relevant distance scale. The inequality is a lower bound because additional coordination costs (negotiation, synchronization, conflict resolution — see Working Notes) contribute beyond the closure-defect mechanism. When the macro-agent’s correction of internal mismatch is less efficient than its correction of external mismatch (because internal mismatches may be harder to diagnose — they look like environmental change from the macro perspective), the actual overhead exceeds this bound.

Consequences (dimensionally correct):

- The realized external tempo $\mathcal{T}_c^{\text{ext}}$ (the share of composite tempo directed at *external* mismatch) chains under the joint baseline — the coordination-overhead definition plus the bridge-lemma lower bound give the middle inequality, and externally-directed tempo cannot exceed total tempo: $\mathcal{T}_c^{\text{ext}} \leq \mathcal{T}_c \leq \mathcal{T}_{\text{joint}} - \frac{\varepsilon^* \nu_c}{\|\delta_{\text{critical}}\|}$ ($= \sum_i \mathcal{T}_i - \varepsilon^* \nu_c / \|\delta_{\text{critical}}\|$ under cross-agent noise independence)
- **Brooks’s Law.** Adding an agent increases $\mathcal{T}_{\text{joint}}$ (by its marginal joint contribution — which cross-agent correlation can make much smaller, or larger, than its solo tempo) but may also increase ε^* . The turning point — where adding a member lowers realized external tempo — occurs when: $\frac{\Delta \varepsilon^* \nu_c}{\|\delta_{\text{critical}}\|} > \Delta \mathcal{T}_i$. The closure-defect increase (expressed as tempo penalty) exceeds the new member’s tempo contribution.
- **Composite persistence condition.** The composite persists iff $\mathcal{T}_c > \rho_{\text{eff}} / \|\delta_{\text{critical}}\|$; a *necessary* condition in joint-baseline terms is therefore $\mathcal{T}_{\text{joint}} > \frac{\rho_{\text{ext}} + \varepsilon^* \nu_c}{\|\delta_{\text{critical}}\|}$ which separates internal-coordination overhead from external challenge while keeping tempo units on the LHS and disturbance-rate units on the RHS (it reduces to the historical $\sum_i \mathcal{T}_i$ form exactly under cross-agent noise independence).

Discussion.

Equality conditions. $\mathcal{T}_c = \mathcal{T}_{\text{joint}}$ (i.e. $C_{\text{coord}} = 0$) requires: 1. **Orthogonal Routing:** The information dependency DAG exactly matches organizational boundaries (no costly cross-talk required). 2. **Perfect Shared Intent:** No tempo is lost to internal negotiation or conflicting pursuit of the macro-objective. 3. **Lossless Fusion:** No joint information in the pooled streams is discarded by the composite’s aggregation.

These are stated as sufficient conditions; whether they are also necessary is open. The further equality $\mathcal{T}_{\text{joint}} = \sum_i \mathcal{T}_i$ is a *separate* condition — cross-agent noise independence — and its two-sided failure is what retired this segment’s historical additive form: under common-source noise a zero-defect, zero-overhead composite still shows $\mathcal{T}_c < \sum_i \mathcal{T}_i$ (the additive baseline is inflated), and under anti-correlated noise a zero-defect composite shows $\mathcal{T}_c > \sum_i \mathcal{T}_i$ (triangulation). The old condition 3 (“no net macro-information loss”) implicitly carried both the fusion condition and the independence condition; they are now separated.

Brooks’s Law. Adding more agents (increasing $\mathcal{T}_{\text{joint}}$ by their marginal joint contribution — which cross-agent correlation can shrink toward zero or amplify) only increases the composite tempo $\mathcal{T}_c$ if the corresponding increase in closure defect ε^* doesn’t let C_{coord} dominate. The model provides a formal analog of Brooks’s Law: if coordination overhead (C_{coord}) grows faster than joint capability ($\mathcal{T}_{\text{joint}}$) when adding agents, then adding people to a late project makes it later — and a new member whose observation noise echoes the existing team’s contributes almost nothing to $\mathcal{T}_{\text{joint}}$ while still adding coordination links, the sharpest Brooks case. Whether this specific mechanism (coordination overhead consuming tempo) is the dominant cause in practice is an empirical question.

Connection to 15.1. The persistence condition for the composite agent requires $\mathcal{T}_c > \rho_{\text{ext}}/\|\delta_{\text{critical}}\|$. Since $\mathcal{T}_c = \mathcal{T}_{\text{joint}} - C_{\text{coord}}$, high coordination overhead can push the composite below the persistence threshold, causing it to disintegrate as a coherent entity — even though each sub-agent individually persists.

Wrapping construction as a Brooks’s-Law instance. Class coercion of a Class 2 (Partial) or Class 3 (Coupled) component via wrapping (`#der-class-coercion-via-wrapping`) is a concrete instance of the same coordination-overhead form. A wrapper makes $K \geq 2$ component calls per macro-step (one for each goal-blind belief-update query, one for the goal-conditioned strategy-update query, more in richer wrapper designs). The wrapper-level macro-update rate is $\nu_W = \nu_A/K$ where ν_A is the underlying component’s nominal call rate, so the wrapper-level tempo is $\mathcal{T}_W \leq \mathcal{T}_A^{\text{nominal}} - C_{\text{coord}}^{\text{wrap}}$, with the coordination overhead $C_{\text{coord}}^{\text{wrap}}$ scaling with K . The cost of class coercion — the structural enforcement of directed separation at the wrapper level when the underlying component does not provide it — is paid in macro-tempo. Strict-wrapping (W_1) regimes have $K \geq 2$ minimum; partial-wrapping (W_2) regimes have $K = 1$ for the LLM call plus some parsing overhead, but the structural-separation guarantee is correspondingly weaker (see `#der-class-coercion-via-wrapping`).

Working Notes.

- **Migration note (2026-05-09 GUC rename):** Class 2 $\leftrightarrow$ Class 3 swap. Pre-2026-05-09: Class 2 = fully merged, Class 3 = partially modular. Post: Class 2 = Partial, Class 3 = Coupled. Removed at candidate stage per FORMAT.md Gate 4.
- The lower bound $C_{\text{coord}} \geq \varepsilon^* \nu_c / \|\delta_{\text{critical}}\|$ captures the minimum tempo-equivalent overhead from closure defect. Additional coordination costs likely contribute beyond this: negotiation costs (reaching shared decisions), synchronization costs (waiting for slower agents), and conflict resolution costs (undoing contradictory actions). These are not captured by ε^* because they arise from the coordination PROCESS, not from the closure DEFECT. A fuller model would add: $C_{\text{coord}} = \varepsilon^* \nu_c / \|\delta_{\text{critical}}\| + C_{\text{negotiation}} + C_{\text{sync}} + C_{\text{conflict}}$, with all terms in tempo units.
- **Heterogeneity drives closure defect.** For agents with equal correction rates ($\alpha_1 = \alpha_2$), $\varepsilon^* = 0$ and $C_{\text{coord}} = 0$. For agents with different correction rates, $\varepsilon^* \propto |\alpha_1 - \alpha_2|$ — heterogeneity drives closure defect, confirmed analytically in the 2-agent toy case.
- The internal/external mismatch distinction is formalized through the bridge lemma: trajectory error e_t is the internal mismatch (divergence between macro-prediction and micro-reality, units of distance), and the tempo spent reducing e_t is the coordination overhead. The steady-state internal mismatch is $\varepsilon^* \nu_c / \alpha_c$ (a distance, since α_c has units $[\text{time}^{-1}]$), and maintaining this at the critical-distance scale requires ongoing tempo expenditure of $\varepsilon^* \nu_c / \|\delta_{\text{critical}}\|$.
- **Norm specification note.** The closure defect ε^* depends on the norm choices in 13.1. For estimation-type agents, the natural norm is Mahalanobis (weighted by inverse prediction-error covariance). In these norms, ε^* measures the ratio of missed cross-correction to expected estimation uncertainty — a quantity with direct tempo interpretation: it’s the fraction of macro-tempo consumed by correcting errors that optimal cross-exploitation would eliminate.
- **Provenance (2026-07-16).** The signed reformulation (joint baseline, super-additive witness, retirement of the unconditional additive bound) was derived in `spikes/spike-composite-tempo-synergy-2026-07-16.md`, discharging the same-day Working-Note flag; the witness numerics are re-runnable from the spike. Do not restore the additive baseline without checking cross-agent noise independence for the application. Audit-gold quotes elsewhere in these notes that reference the retired additive form are first-encounter archaeology, kept verbatim.

- **Dynamic- θ regime (from the confirmer $\neq$ author verification, 2026-07-16).** Under random-walk θ with steady-state Kalman sub-agents, the *optimal-gain* fused composite no longer commutes exactly ($\epsilon^* > 0$ measured when $\eta_c \neq \eta_i$); the commuting composite (running the sub-agents' gain on the averaged channel) still beats $\sum_i \mathcal{T}_i$ for $\rho < -1/2$ in the own-steady-state reading. A fidelity-tempo tradeoff therefore *does* appear in the dynamic regime even though the static witness has none — worth its own small derivation if a consumer needs the dynamic threshold sharpened.
- **(A4) uniformity wrinkle (static witness).** The running-mean gain decays ($\eta_m = 1/(m+1)$), so no uniform $\alpha_c > 0$ holds over an infinite horizon in the static case; the sub-agents share the wrinkle symmetrically (their gains decay identically), so it cancels from the ratio, and the dynamic- θ steady state repairs it outright.
- **(P1) holds with $\epsilon_I = 0$ exactly** for equal variances — the channel average is exactly sufficient for θ within the pooled pair, so the witness's projection loses *no* predictive information (stronger than admissibility requires); 12.2 is met via the shared-objective route (G_c, π_c trivial in the passive-estimation instance).
- **ϵ^* scaling with N — resolved (no longer open).** The poly-vs-superlinear framing was an accumulation-type confound (a per-step residue asked an accumulation question); re-typed it dissolves — see 13.1 (the N -agent-scaling and strategy-DAG-projection rows + Discussion): $\epsilon^*(N)$ is dimension-free-zero in the benign linear-Gaussian-stationary regime (all N , all coupling), graph-Laplacian-bounded with no exponential regime under compression, and order-incompatibility-invariant ($\leq |S| \log 2$, N -free) for strategy-DAG composition. Brooks's-Law collapse, if it occurs, is driven by coordination overhead C_{coord} , not by $\epsilon^*(N)$ growth (the Brooks's-Law-floor corollary is in 8.3 §"Composing heterogeneous strategy DAGs"); CHANGELOG 2026-05-19.
- **Track E surface-back of catalog citations (2026-05-22).** Three prior-art neighbors surfaced 2026-05-22 from Track E catalog at ref/prior-art-analysis/08-composite-agency.md (Pillar 4 + key anchor papers), each marked [^cat-2026-05-22] for verification-deferred attribution:
 - **Janow 2009, *A Fundamental Limit on Productivity in Organizations: Collaborative Entropy Costs*^a** — closest single-paper precedent for the tempo-tax claim. First-principles Shannon-entropy model of coordination cost as productivity cap. Uses Shannon entropy rather than persistence-bound cybernetic accounting; AAT's contribution is the integration with #form-composition-closure's ϵ^* machinery and the persistence-framework dimensional accounting (Composite Tempo Inequality, now in its joint-baseline form $\mathcal{T}_c = \mathcal{T}_{\text{joint}} - C_{\text{coord}}$).
 - **Gurvich & Van Mieghem 2015, *Collaboration and Multitasking in Networks: Architectures, Bottlenecks, and Capacity*^b** — strongest nearby formal result on synchronization-induced bottleneck idleness as unavoidable coordination overhead. Structural cousin to C_{coord} — same phenomenon under queuing-theoretic vocabulary rather than persistence-framework vocabulary.
 - **Bamieh, Jovanović, Mitra & Patterson 2011, *Coherence in Large-Scale Networks: Dimension-Dependent Limitations of Local Feedback*^c** — large networks can lose macro-coherence as they scale even when local regulation is good. Structural cousin to Brooks's-Law: dimension-dependent loss of coherence as composition order grows.

Primary-source verification queued in the BG2 cluster — see #disc-identifiability-floor Working Notes for the verification-targets list. Catalog citations had prior Pillar-style search support but were not verified by the current executor at landing time.

Incidental audit gold (lift 2026-05-30). Cross-audit “wandering thoughts” / §14-ideation harvested from the de-novo auditors' working dirs, deduplicated across substrates and lightly attributed. Orthogonal pedagogical / framing material, kept separate from the certified theory-fix findings (handled elsewhere — the F116–F121 tempo-double-counting cluster is routed to the findings track; see report flags). **Coverage:** four dirs reached a digested reflection (526815, 773921, 829314, 849201). Universally high felt-value — the Brooks's-Law derivation and the dimensional accounting were the converged standouts. Substrate attribution inferred from voice where not explicit.

1. Candidate Brief prose / pre-prose.

- **Brooks's Law as a strict thermodynamic inequality.** The converged star of this segment: Fred Brooks observed “adding manpower to a late software project makes it later” and attributed it to $O(N^2)$ communication overhead; AAT gives the *exact* inequality for when the catastrophic slowdown actually happens — $\Delta\epsilon^* \nu_c / \|\delta_{\text{critical}}\| > \Delta\mathcal{T}_i$ — translating “a famous software engineering aphorism into a strict thermodynamic inequality” (Gemini, AUDIT-WORKING-829314 — “an absolute triumph”; Claude, AUDIT-WORKING-849201; Claude, AUDIT-WORKING-773921 — “a formal, computable physics equation for team size limits”). The worked mechanism: a junior engineer's small $\Delta\mathcal{T}_i$ is exceeded by the C_{coord} the seniors spend correcting the internal mismatch the addition created (explaining the codebase, fixing bad PRs, unblocking dependencies). A strong Brief/Discussion anchor.
- **The thermodynamic equation of teamwork.** $C_{\text{coord}} = \sum_i \mathcal{T}_i - \mathcal{T}_c$ glossed as a conservation law: “energy in (sum of tempos) equals useful work ($\mathcal{T}_c^{\text{ext}}$) plus heat (C_{coord})” (Claude, AUDIT-WORKING-849201). Makes the sub-additive bound legible as a first-law accounting.

2. Candidate Discussion.

- **The cost of diversity — heterogeneity drives closure defect.** The existing “ $\epsilon^* \propto |\alpha_1 - \alpha_2|$ ” Working Note, pedagogically unpacked: teams of clones / homogeneous cultures stay in sync “without talking” ($\epsilon^* \approx 0$); a fast learner paired with a slow learner constantly generates internal mismatch for the slow one (by unexpectedly changing APIs / processes), consuming coordination overhead to re-align — so “cognitive diversity in a team mathematically generates coordination overhead.” The honest two-sidedness: the *benefit* of diversity (wider structural reserve R , more accurate M_t) must outweigh this strict ϵ^* penalty for a heterogeneous team to be viable — “AAT provides the calculus to weigh that tradeoff objectively” (Gemini, AUDIT-WORKING-829314; Claude, AUDIT-WORKING-773921 — “if a model assumes diverse teams coordinate for free, it violates this math”). Candidate Discussion expansion of the heterogeneity Working Note. (*Keep it two-sided — the deflationary “diversity is just cost” reading would be a misframe; the point is a measurable tradeoff, not a verdict.*)
- **C_{coord} is a floor, so Brooks's Law is worse than the equation states.** The closure-defect term captures only the overhead of imperfect macro-representation; negotiation, synchronization, and conflict-resolution costs add more on top — so the inequality is “the absolute theoretical minimum overhead for simply representing the team as a single unit” (Claude, AUDIT-WORKING-773921). Already an existing Working Note ($C_{\text{coord}} = \epsilon^* \nu_c / \|\delta_{\text{critical}}\| + C_{\text{negotiation}} + C_{\text{sync}} + C_{\text{conflict}}$); converged-on framing for why the bound is conservative-in-the-honest-direction.

3. Follow-up items.

- **Wrapping construction as a Brooks's-Law instance (cross-volume).** An LLM wrapper framework that runs K API calls per macro-step to verify/plan is exactly a multi-agent composition with high coordination overhead: “you burn micro-tempo (ν_A/K) to buy macro-level structural separation” (Claude, AUDIT-WORKING-773921). Connects this segment's tempo accounting to 13.4 / 13.5's tempo-cost claim — a candidate explicit cross-reference (the in-composition segment already cites this Brooks's-Law form).

4. Readers often ask / wonder.

- **“Does the convergence with established results (Janow, Bamieh, Gurvich–Van Mieghem) confirm the derivation?”** Auditors read the Track-E surface-back literature (collaborative-entropy productivity caps; dimension-dependent loss of macro-coherence in large networks; synchronization-bottleneck idleness) as evidence that “AAT's first-principles derivations are converging perfectly with established theorems in operations research and distributed control” (Claude, AUDIT-WORKING-773921). A convergence-as-coherence-evidence signal; the prior-art neighbors already live in Working Notes (verification-deferred). Worth a Discussion sentence framing the convergence honestly once those citations are verified.

5. Candidate figures.

- **Tempo flow with a coordination-overhead split, plus a double-count warning.** Aggregate sub-agent tempo $\sum_i \mathcal{T}_i$ flowing toward external correction, with closure defect ε^* diverting part of it into coordination C_{coord} . The diagram explicitly carries a *warning split*: count closure either as an added disturbance burden ($\rho_{\text{eff}} = \rho_{\text{ext}} + \varepsilon^* \nu_c$) or as tempo overhead, “but do not subtract it twice” (Codex/Claude, AUDIT-WORKING-526815, “diagram attempt”). The double-count warning the figure encodes is also the substance of the F116–F119 correctness findings flagged in the report — the figure and the finding point at the same ledger discipline.

^aCitation surfaced 2026-05-22 from the Track E catalog at ref/prior-art-analysis/ (intermediate work artifacts that captured Pillar-style prior-art searches). Catalog has more verification support than raw Undermind synthesis but less than full primary-source reading. Verification queued with the BG2 cluster.

^bCitation surfaced 2026-05-22 from the Track E catalog at ref/prior-art-analysis/ (intermediate work artifacts that captured Pillar-style prior-art searches). Catalog has more verification support than raw Undermind synthesis but less than full primary-source reading. Verification queued with the BG2 cluster.

^cCitation surfaced 2026-05-22 from the Track E catalog at ref/prior-art-analysis/ (intermediate work artifacts that captured Pillar-style prior-art searches). Catalog has more verification support than raw Undermind synthesis but less than full primary-source reading. Verification queued with the BG2 cluster.

13.3 Hypothesis: Directed Separation Under Composition

The question whether the composite-level directed-separation condition holds: if every sub-agent satisfies 6.3 individually, does the composite macro-agent also satisfy it? The answer is *conditional on the composite’s routing structure* (12.1).

The framework draws a careful distinction first. Sub-agents’ goal-driven actions shape the environment, which other sub-agents observe. This is **not** a violation of directed separation — it is the same mechanism the single-agent result explicitly allows: goals influence *events* through action, but processing of realized events is goal-blind. Agent *B* observing environmental changes caused by Agent *A*’s goal-driven behavior is *B* processing a realized event goal-blindly. The observations carry information about *A*’s goals, but that is a property of the event’s content, not a failure of goal-blind processing. Similarly, individual messages reflecting individual goals is action through policy, not a routing-structure dependence on the composite’s goals.

The substantive question is whether the **routing structure** depends on the composite’s goal. Two cases. **Case 1 — goal-blind routing**: when neither communication topology nor protocol depends on the composite’s goals (military command structures with doctrinal communication protocols; software teams with defined code-review processes; multi-agent AI systems with protocol-specified message passing), directed separation *survives* at the composite level. **Case 2 — goal-dependent routing**: when the topology changes based on objective (different reporting chains activated depending on mission) or the protocol changes (different intelligence products shared depending on objective), the composite’s effective observation function acquires a goal argument. Even if each sub-agent processes its observations goal-blindly, *the set of observations reaching each sub-agent depends on the composite’s goal through the routing function*. Directed separation *fails* at the composite level.

This refinement establishes that **composite-level class membership depends on routing structure, not just sub-agent class**. A composite of Class-1-(Separated) sub-agents with goal-blind routing remains Class 1 at the composite level. A composite of Class-1-(Separated) sub-agents with goal-*dependent* routing becomes effectively Class 3 (Coupled) at the composite level — and the coupled-formulation machinery applies at the macro level even though every sub-agent individually satisfies directed separation. The architectural classification is therefore *not preserved under composition with goal-dependent routing*. For the special case of a single-component composite (wrapper-around-component), 13.4 promotes the equivalent result from hypothesis to derived; the general *N*-agent case treated here remains a hypothesis.

Formal Expression.

Consider N agents $A_1, \dots, A_N$, each satisfying directed separation individually: $M_{i,\tau+} = f_M^{(i)}(M_{i,\tau-}, e_{i,\tau})$ with no $G_t^{(i)}$ argument.

Hypothesis (directed-separation-under-composition, extending directed-separation to composites).

The question. Directed separation is about **processing**, not **selection** (6.3, scope condition). A single agent's goals affect which events it seeks (through π), but f_M processes whatever event arrives without reference to G_t . At the composite level, the analogous question is: does the composite's routing structure R_t (12.1) depend on the composite's goals G_t^c ?

Note: sub-agents' goal-driven actions shape the environment, which other sub-agents observe. This is NOT a violation of directed separation — it is the same mechanism directed separation explicitly allows at the single-agent level: goals influence events through action, but processing of realized events is goal-blind. Agent B observing environmental changes caused by A 's goal-driven behavior is B processing a realized event goal-blindly. The observations carry information about A 's goals, but that is a property of the event's content, not a failure of goal-blind processing. Similarly, individual messages reflecting individual agents' goals is action through policy, not a routing-structure dependence on G_t^c .

Case 1: Goal-blind routing.

If the routing structure satisfies $R_t \perp G_t^c$ (12.1, goal-blind routing) — neither the communication topology $\mathcal{N}_t$ nor the protocol $c_t^{(j \rightarrow i)}$ depends on the composite's goals — then:

Hypothesis (Case 1).

- Each sub-agent processes observations goal-blindly (individual directed separation)
- The routing is goal-blind (by construction)
- Therefore f_M^c is G_t^c -independent

Directed separation **survives** at the composite level. Examples: military command structures with doctrinal communication protocols, software teams with defined code-review processes, multi-agent AI systems with protocol-specified message passing.

Case 2: Goal-dependent routing.

If R_t depends on G_t^c — either the topology $\mathcal{N}_t$ changes (different reporting chains activated depending on the mission) or the protocol $c_t^{(j \rightarrow i)}$ changes (different intelligence products shared depending on the objective) — then the composite's effective observation function has a goal argument:

Hypothesis (Case 2).

$$o_c = h^c(\Omega, a_{\text{micro}}, G_t^c, \xi) \quad (13.26)$$

Even if each sub-agent's $f_M^{(i)}$ processes o_i goal-blindly, the **set of observations reaching each sub-agent** depends on G_t^c through the routing function. Directed separation **fails** at the composite level. The composite is analogous to a Class 3 (Coupled) architecture: goal content shapes the information pathway, not through individual interpretation but through collective routing.

Epistemic Status.

Conditional on the routing structure of the composite. Max attainable: conditional (the two cases are genuine architectural alternatives). Previously discussion-grade; upgraded after routing formalization in 12.1 and architectural classification promotion in 6.3 (2026-04-01).

Foundation status.

1. The **architectural classification** (Class 1/2/3) is in 6.3's Formal Expression with formal operationalization ($\kappa_{\text{processing}}$). Status: robust qualitative.
2. The **routing structure** R_t is defined in 12.1 with a formal goal-independence condition ($R_t \perp G_t^c$). The definition decomposes into topology independence and protocol independence. Whether this captures all relevant ways that composite information flow can depend on goals is an open question.
3. The **admissible coarse-graining** Λ from 13.1 now has specified admissibility constraints: (A1)-(A4) for macrodynamics and (P1)-(P3) for projections. The remaining gaps are validation (exercising the machinery on a purposeful multi-agent example) and computability (P1 requires conditional mutual information, tractable only for linear-Gaussian), not missing definitions.

The logic of each case is sound given the routing definition. Case 1: goal-blind routing + goal-blind processing = goal-blind composite. Case 2: goal-dependent routing means the composite's observation function depends on G_t^c , regardless of individual processing. Both follow from directed separation's scope condition.

The claim that the architectural classification lifts to composition is **structurally motivated** by 11 (the theory must give consistent answers at every level of description). This is an argument from theoretical coherence, not a derivation.

Discussion.

What this is and what it isn't. This segment provides a two-case taxonomy for directed separation under composition. It is not yet a derivation — the remaining caveat (admissibility placeholders) prevents that. It IS a useful classification that identifies which composites fall within Part III's scope (Case 1) and which require a coupled formulation (Case 2).

Constructive special case: wrapper-around-component. For the special case of a single-component composite — one underlying primitive component embedded inside an external scaffold — the question of when directed separation holds is *derived* by #der-class-coercion-via-wrapping. Under explicit conditions on the component (admissibility of goal-blind queries, stationary conditional, no implicit goal-inference) and on the wrapper's update maps (type signatures that exclude G_W from the belief-update path), directed separation holds at the wrapper level by structural commitment. The construction promotes the wrapper-around-component case from hypothesis to derived; the general N -agent composition case (Case 1 vs. Case 2 above) remains a hypothesis pending further work.

Most composites of interest are Case 1. Fixed communication structures are the norm in designed multi-agent systems: military doctrine specifies information flow regardless of the current mission; software development processes define code review, CI, and deployment pipelines independent of the feature being built; multi-agent AI protocols specify message formats and channels. Goal-dependent routing (Case 2) is rarer: crisis management (where communication structure changes with threat type), ad hoc teams assembled for specific objectives, and attention-based multi-agent architectures where query routing depends on the shared goal.

Connection to logogenic agents. LLM-based agents are individually Class 3 (Coupled; goal-conditioned attention). Whether composites of LLM agents are Case 1 or Case 2 depends on whether the inter-agent communication protocol is fixed or goal-dependent. A multi-agent LLM system with fixed API contracts between agents is Case 1 at the composite level even though each agent is individually Class 3 (Coupled). The internal architecture of each agent and the composition architecture are separate questions.

Adversarial implications. Goal-dependent routing (Case 2) makes the composite’s objective partially observable through its communication patterns. An adversary who can observe which information is being routed where can infer the composite’s current goal. This creates a meta-strategic incentive for fixed routing: it preserves not just epistemic hygiene but operational security.

Working Notes.

- **Migration note (2026-05-09 GUC rename):** Class 2 $\leftrightarrow$ Class 3 swap. Pre-2026-05-09: Class 2 = fully merged, Class 3 = partially modular. Post: Class 2 = Partial, Class 3 = Coupled. Removed at candidate stage per FORMAT.md Gate 4.
- **Goal-information leakage is a separate phenomenon.** In any composite, sub-agents’ goal-driven actions shape the environment, so observations carry statistical information about goals. This is real and may matter for adversarial dynamics, trust calibration, and OPSEC — but it is NOT a directed-separation issue. It’s the normal action-environment coupling that directed separation explicitly allows. If this phenomenon deserves formalization (and it may — the mutual information $I(o_c; G_t^c \mid \Omega_t)$ is well-defined and operationally relevant), it should be its own segment about goal-information leakage, not a case within directed-separation analysis. An earlier draft of this segment conflated the two; this was caught by external review.
- **Partial goal-dependence.** Real composites may have mostly-fixed routing with occasional goal-dependent exceptions (a military unit that follows doctrinal comms but occasionally switches to crisis-specific channels). The error from treating such composites as Case 1 depends on the frequency and magnitude of the goal-dependent exceptions. A formal treatment would need to quantify this, possibly through a “routing independence fraction.”
- **Testable prediction.** Organizations that restructure communication channels based on the current strategic objective should exhibit lower prediction accuracy at the organizational level (the shared model becomes goal-contaminated through routing, not just through environmental coupling). Compare prediction accuracy before and after mission-dependent communication restructuring.
- **Composite-class-inheritance table (canonical home: #der-directed-separation; added 2026-05-21 per Phase 5 cross-segment ripple).** This hypothesis segment’s Case 1 / Case 2 dichotomy was the original framing for composite-level class lift. The 2026-05-21 strategic-composition cycle (spikes/strategic-composition-class-3-attempt-2026-05-21/) extended the analysis along a third axis (substrate sharing with G^c -allocation, per 01-STRENGTHEN-ATTEMPTS.md §5 (R4)) and produced the *axis-decomposed composite-class-inheritance table* now living in #der-directed-separation §“Composite-level class inheritance.” That table is the consolidated home for the sub-agent-class $\times$ routing $\times$ substrate $\rightarrow$ composite-class lookup; this segment’s Case 1 / Case 2 dichotomy remains the underlying hypothesis (still hypothesis-grade — the N -agent general case is not derived; the wrapper-around-component special case is derived in #der-class-coercion-via-wrapping). The dynamic-regime axis (#disc-dynamic-regime-axis) — goal alignment vs strategic interaction — is *independent* of the routing-structure axis this segment treats, and a clean separation: alignment work moves composites along the regime axis; routing-structure work moves them along the architectural-class axis.

Incidental audit gold (lift 2026-05-30). Cross-audit “wandering thoughts” / §14-ideation harvested from the de-novo auditors’ working dirs, deduplicated across substrates and lightly attributed. Orthogonal pedagogical / framing material, kept separate from the certified theory-fix findings (handled elsewhere — the F122–F126 cluster, esp. the Case-1-needs-goal-blind-projection-too point, is routed to the findings track; see report flags). **Coverage:** two dirs reached a digested reflection (526815, 849201). Lower note-count, but high felt-value on the OPSEC and LLM-composite-inversion observations. Substrate attribution inferred from voice where not explicit.

1. Candidate Brief prose / pre-prose.

- The Case 1 / Case 2 line stated plainly as an infrastructure test: a composite stays Class 1 (Separated) when “the infrastructure doesn’t change based on the goal” ($R_t \perp G_t^c$) and each sub-agent’s f_M is goal-blind; it slides to Class 2 when “the infrastructure shifts based on the goal (e.g., activating crisis channels), [so] the composite’s observation function becomes goal-dependent” (Claude, AUDIT-WORKING-849201). A compact Brief gloss of the dichotomy.

2. Candidate Discussion.

- **Goal-dependent routing is an OPSEC vulnerability — derived from the structure.** The standout insight: when routing depends on the composite goal, “adversaries can infer the goal just by watching the routing topology” — so Case-2 routing is not only an epistemic-hygiene failure but a security leak, “a brilliant practical insight derived directly from the theoretical structure” (Claude, AUDIT-WORKING-849201). A candidate Discussion addition turning the routing distinction into an operational consequence (and it resonates with the consciousness-infrastructure goal-blind-routing imperative lifted into 12.1).
- **The LLM-composite inversion.** “While individual LLMs are Class 2/3, a multi-LLM system with fixed API contracts might be Case 1 at the macro level” — a “fascinating inversion of expectations” where composition *recovers* separation that the components individually lack (Claude, AUDIT-WORKING-849201). The segment treats this; the framing is a candidate reader-orienting Discussion line. (*Correctness caveat flagged in the report: F124 notes this belongs under the wrapper-derived special case 13.4, since the general Case-1 setup assumes each sub-agent is individually separated — keep the tier/scope honest if promoted.*)

3. Follow-up items.

- **Case 1 needs goal-blind projection, not just goal-blind sub-agents + routing.** A real structural sharpening (routed as a finding, noted here for completeness): the Case-1 conclusion also requires the macro projection / coarse-graining and macro-update interface to be goal-blind — a goal-conditioned Λ , aggregation window, or macro-observation definition can reintroduce coupling even when sub-agent processing and routing are clean. The proposed sharper theorem shape: “goal-blind sub-agent processing, goal-blind routing, *and* goal-blind projection/update together preserve composite-level directed separation” (Codex/Claude, AUDIT-WORKING-526815, F122). This is a strengthen-the-hypothesis candidate, not framing gold — see report off-ramp flag.

4. Readers often ask / wonder.

- “Isn’t a sub-agent inferring another’s goal from its actions already a separation failure?” No — and the segment’s existing “goal-information leakage” Working Note resolves exactly this trap: inferring goals from observed actions is normal, valid Bayesian inference (the action-environment coupling directed separation explicitly allows), distinct from a *structural* flaw in the processing pipeline. A fresh reader independently flagged this distinction as “resolving a major conceptual trap” (Claude, AUDIT-WORKING-849201). Strong signal the distinction is a natural reader stumble worth keeping prominent.

Belongs elsewhere.

- (*No distinct belongs-elsewhere gold beyond the cross-links already noted — the OPSEC and goal-blind-routing-imperative material co-locates naturally with 12.1 and the consciousness-infrastructure reach lifted there.*)

13.4 Derived: Class Coercion via Wrapping

A *constructive* result of substantial practical importance. A Class 2 (Partial) or Class 3 (Coupled) component (one whose forward pass entangles belief-update and goal-conditioning) can be embedded inside an external scaffold whose state $X_W = (M_W, G_W)$ is updated by *structurally distinct query channels*: **goal-blind queries** to the component update M_W ; **goal-conditioned queries** update G_W . Under stated conditions on the component, directed separation holds at the wrapper level *by construction*, and the composite system is Class 1 (Separated) per #der-directed-separation — even though the underlying component is not. This is the constructive direction of #hyp-directed-separation-under-composition for the wrapper-around-component special case: a procedure for *making* directed separation hold when the underlying component does not provide it.

The wrapper has four type-signed components: a *belief-side query selector* that chooses the model-update query from belief and observation only — *no goal argument*; a *strategy-side query selector* that may depend on the goal; a *belief update map* that updates M_W from prior belief, observation, the query made, and the component’s response — *no goal*

argument; and a *strategy-update map* that may depend on the goal. The wrapper makes at least two component calls per macro-step: one goal-blind for the model update and one goal-conditioned for the purposeful-state update.

The result requires conditions on the component. **(C1) Goal-blind admissibility** — the component admits non-trivial goal-blind queries; the framework partitions components into *Class A* (goal-blind by design: POMDP belief-state filters, world models, sensory pipelines, retrieval systems), *Class B* (admit a goal-blind query mode alongside goal-conditioned ones: LLMs in summarization/fact-extraction modes, hybrid RL with separable value/policy), and *Class C* (fundamentally goal-conditioned: pure end-to-end goal-conditioned policy networks), with the construction applying to Classes A and B. **(C2) Stationary component conditional** — the component’s output distribution conditional on input is fixed during operation (adaptation-during-deployment systems are out of scope). **(C2’) No goal-correlated cross-call state** — the component carries no information about the latent operator goal across the goal-blind / goal-conditioned call boundary; this is the condition on which the *structural* W_1 leakage bound depends. **(C3) No implicit goal-inference** — the component’s response to a goal-blind query does not depend on G_W via inference from query patterns; for pretrained components like LLMs, (C3) holds *exactly* only when query content is statistically independent of G_W in the pretraining distribution. The framework supplies **two theorems**: an *exact* form under (C1)–(C3), and an *approximate* form under (C1)–(C2) plus a KL-leakage bound, where the wrapper-level leakage on M_W updates is bounded by the same bound (small leakage in, small leakage out via the data-processing inequality).

The construction supports three **wrapping regimes** of decreasing strictness, distinguished by where structural separation lives. **W_1 (strict wrapping)** uses separate q_M and q_G calls per macro-step; separation lives at the *query boundary* and the residual leakage is the *selection-channel* quantity $\kappa_{W_1}^{\text{sel}} = I(A(q_M); G^{\text{op}}) \leq I(q_M; G^{\text{op}})$ — bounded *structurally* by the goal-content of the wrapper’s query-selection policy (under (C1)+(C2’)). **W_2 (partial wrapping)** uses one goal-conditioned call per macro-step with a typed parsed response routing updates to M_W vs G_W slots; separation lives at the *write boundary* and leakage is bounded only *behaviorally* — by the component’s compliance with the prompted instruction-to-separate. **W_0 (no wrapping)** runs the raw Class 2 / Class 3 component with no separation commitment. What differs across the regimes is *what determines* the leakage rate — and, for W_1 , whether the structural bound is available at all: it requires (C2’), failing which W_1 too degrades to a behavioral bound (`#disc-w1-structural-bound-boundary`). The hierarchy refines the Class 1 (Separated) cell of `#der-directed-separation` with a *Class-1-by-structure* (W_1 or natively goal-blind) vs *Class-1-by-behavior* (W_2) sub-distinction.

The result is *load-bearing* for the framework’s treatment of LLM agents: an LLM is internally Class 3 (Coupled), but an LLM-agent *system* (LLM + tools + memory + monitoring) can be designed with modular topology that recovers Class 1 status at the system level. The construction’s cost is paid in two places — *more component calls per macro-step* (the Brooks’s-Law tempo overhead derived in the companion segment) and a *residual leakage rate* bounded structurally under W_1 or behaviorally under W_2 . The companion segment `#der-class-coercion-in-composition` establishes that the wrapped system is also a valid AAT composite agent (satisfying (A1)–(A4) of `#form-composition-closure`) and inherits the sector-persistence template at the wrapper level.

Formal Expression.

Setup. Let $A : \mathcal{I}_A \rightarrow \mathcal{O}_A$ be a primitive component, treated by the wrapper as a black-box oracle: the wrapper issues queries (inputs) and consumes responses (outputs), without access to A ’s internal state. $\mathcal{Q}_A \subseteq \mathcal{I}_A$ is the set of admissible queries.

A **wrapper** W over A has state $X_W = (M_W, G_W) \in \mathcal{X}_M \times \mathcal{X}_G$ with $\mathcal{X}_G = \mathcal{X}_O \times \mathcal{X}_\Sigma$ per `#def-strategy-dimension`. The wrapper interacts with an environment via observations $o_W \in \mathcal{O}_W$ and actions $a_W \in \mathcal{A}_W$.

[*Definition (wrapper-update-maps)*] The wrapper’s update at macro-step m uses four type-signed components:

- **Belief-side query selector:** $q_M : \mathcal{X}_M \times \mathcal{O}_W \rightarrow \mathcal{Q}_A$. The wrapper chooses the query for M_W updates from belief and observation only — *no G_W argument*.
- **Strategy-side query selector:** $q_G : \mathcal{X}_M \times \mathcal{X}_G \rightarrow \mathcal{Q}_A$. May depend on G_W .

- **Belief-update map:** $f_M : \mathcal{X}_M \times \mathcal{O}_W \times \mathcal{Q}_A \times \mathcal{O}_A \rightarrow \mathcal{X}_M$. Updates M_W from prior belief, observation, the query made, and the component's response. *No G_W argument.*
- **Strategy-update map:** $f_G : \mathcal{X}_G \times \mathcal{X}_M \times \mathcal{Q}_A \times \mathcal{O}_A \rightarrow \mathcal{X}_G$. May depend on G_W .

The external policy $\pi_W : \mathcal{X}_W \rightarrow \mathcal{A}_W$ selects the wrapper's external action.

A macro-step proceeds: construct $q_M(M_W, o_W) \rightarrow$ query $A \rightarrow$ apply f_M ; construct $q_G(M'_W, G_W) \rightarrow$ query $A \rightarrow$ apply f_G ; emit $\pi_W(X'_W)$. The wrapper makes $K \geq 2$ component calls per macro-step in this minimal form (more in richer wrapper designs).

Conditions. [*Conditions (component-admissibility)*] The theorem applies under three conditions on the component A :

(C1) Goal-blind admissibility. $\mathcal{Q}_A$ contains queries whose specification can be constructed from (M_W, o_W) alone — i.e., a non-trivial q_M exists. Components partition into three classes: - **Class A (goal-blind by design).** A 's interface is goal-blind by construction — POMDP belief-state filters, world models, sensory pipelines, retrieval systems, calculators. (C1) holds trivially. - **Class B (admit a goal-blind query mode).** A supports goal-conditioned queries but also goal-blind ones. Large language models in summarization or fact-extraction modes; hybrid RL agents with separable value/policy; multi-modal models. (C1) holds operationally — the wrapper *chooses* to use the goal-blind mode. - **Class C (fundamentally goal-conditioned).** A 's only operating mode requires goal-conditioning. Pure end-to-end goal-conditioned policy networks. (C1) fails; the construction does not apply.

(C2) Stationary component conditional. A 's output distribution conditional on input is fixed during the wrapper's operation: $P(A(\cdot) \mid q)$ does not depend on prior queries or on side information beyond q . Adaptation-during-deployment systems are out of scope.

(C2') No goal-correlated cross-call state. The component's hidden state does not carry information about the latent operator goal G^{op} across the boundary between the goal-blind (q_M) and goal-conditioned (q_G) calls of a macro-step. Equivalently, A 's response to q_M is conditionally independent of G^{op} given q_M and the component's *pre-call* state, and that pre-call state is itself goal-uncorrelated — either reset across the call boundary, or stripped of Σ - and G -content. (C2') is the condition on which the *structural* W_1 leakage bound below depends, and it strengthens (C2): the breaking condition for the structural bound is not online weight-adaptation but mere goal-correlated state persistence across the call boundary — which a frozen-weights LLM with a conversation cache exhibits with no weight adaptation at all. When (C2') fails the structural bound is unavailable and only a behavioral bound remains; this boundary is derived in `#disc-w1-structural-bound-boundary`.

(C3) No implicit goal-inference. A 's response to a goal-blind query does not depend on G_W via inference from query patterns:

$$P(A(q_M) \mid q_M, G_W) = P(A(q_M) \mid q_M) \quad \forall q_M, G_W \quad (13.27)$$

For pretrained components (notably LLMs), (C3) holds *exactly* only when query content is statistically independent of G_W in the pretraining distribution. The approximate form weakens (C3) to a leakage bound (Theorem 2 below).

Theorem 1: Directed separation at the wrapper level (exact form).

Under (C1)–(C3), directed separation holds *exactly* at the wrapper level:

Derived (directed-separation-at-wrapper-exact, from C1+C2+C3).

$$P(M_{W,m+1} \mid M_{W,m}, o_{W,m+1}, G_{W,m}) = P(M_{W,m+1} \mid M_{W,m}, o_{W,m+1}) \quad (13.28)$$

Therefore W is a Class 1 (Separated) architecture per $\#der$ -directed-separation.

Proof. Identify all paths from $G_{W,m}$ to $M_{W,m+1}$ given $(M_{W,m}, o_{W,m+1})$. The update is

$$M_{W,m+1} = f_M(M_{W,m}, o_{W,m+1}, q_M(M_{W,m}, o_{W,m+1}), A(q_M(M_{W,m}, o_{W,m+1}))) \quad (13.29)$$

f_M has no G_W argument by type signature (D-pathway-1 closed). q_M has no G_W argument by type signature (D-pathway-2 closed). The remaining pathway is $A(q_M)$ depending on G_W given q_M . Under (C3), $P(A(q_M) \mid q_M, G_W) = P(A(q_M) \mid q_M)$ — the response is conditionally independent of G_W given q_M . Since q_M is itself a deterministic function of $(M_{W,m}, o_{W,m+1})$, conditioning on $(M_{W,m}, o_{W,m+1})$ determines q_M , and the integrand $P(M_{W,m+1} \mid M_{W,m}, o_{W,m+1}, q_M, A(q_M)) \cdot P(A(q_M) \mid q_M, G_W)$ no longer depends on G_W . The conditional distribution of $M_{W,m+1}$ given $(M_{W,m}, o_{W,m+1}, G_{W,m})$ equals that given $(M_{W,m}, o_{W,m+1})$. ■

Theorem 2: Directed separation (approximate form, C3 weakened to leakage bound).

If (C3) is replaced by a KL-leakage bound

Derived (directed-separation-at-wrapper-approximate, from C1+C2+leakage-bound).

$$D_{\text{KL}}(P(A(q_M) \mid q_M, G_W) \parallel P(A(q_M) \mid q_M)) \leq \kappa \quad \forall q_M, G_W \quad (13.30)$$

then the wrapper-level KL-divergence on M_W updates is bounded by the same κ :

$$D_{\text{KL}}(P(M_{W,m+1} \mid M_{W,m}, o_{W,m+1}, G_{W,m}) \parallel P(M_{W,m+1} \mid M_{W,m}, o_{W,m+1})) \leq \kappa \quad (13.31)$$

The wrapper is *almost-Class-1 (Separated)* with leakage rate $\leq \kappa$. *Proof.* The wrapper-level M_W update is a deterministic function of the component response given the wrapper's other inputs; the data-processing inequality propagates the KL bound from response distribution to wrapper-state distribution. ■

The W_1 structural leakage bound: a selection-channel quantity. The residual goal-leakage that survives in the W_1 regime is a *selection-channel* quantity, not a processing-channel one. The distinction is between two senses of the goal variable: the wrapper's internal purposeful-state register G_W — which q_M structurally does not take as an argument and which Theorem 1 closes by type signature — and the latent *operator goal* G^{op} that generated the situation the wrapper is in, which a competent component can infer from the *content* of the query the wrapper chose. The processing channel ($G_W \rightarrow A(q_M) \mid q_M$) is shut structurally and exactly by Theorem 1; the entire residual leakage is the selection channel ($G^{\text{op}} \rightarrow \text{wrapper history} \rightarrow \text{choice of } q_M \rightarrow A(q_M)$), and its magnitude is governed by the goal-content carried in the query the wrapper selects.

Under (C1) and (C2'), the goal-information reaching the wrapper's belief update through the goal-blind channel is bounded by the goal-content of the wrapper's query-selection policy:

Derived (*W1-selection-leakage-bound, from C1+C2', data-processing inequality*).

$$\kappa_{W_1}^{\text{sel}} := I(A(q_M); G^{\text{op}}) \leq I(q_M; G^{\text{op}}), \quad (13.32)$$

where G^{op} is the latent operator goal and q_M is the (random, history-dependent) goal-blind query drawn under the wrapper's operating policy.

Derivation. Under (C2'), $A(q_M)$ depends on G^{op} only through q_M — there is no goal-correlated cross-call state through which the goal could reach the response while bypassing the query — so $G^{\text{op}} \rightarrow q_M \rightarrow A(q_M)$ is a Markov chain. The data-processing inequality along this chain gives $I(A(q_M); G^{\text{op}}) \leq I(q_M; G^{\text{op}})$. ■

The bound is a property of the wrapper's q_M -selection policy, not of the component's input–output law, and is therefore design-controllable: maximally goal-blind queries (current observation only, no history, no system prompt) drive $I(q_M; G^{\text{op}}) \rightarrow 0$ and recover exact directed separation on the merits; richer, history-laden queries raise $I(q_M; G^{\text{op}})$ and with it the leakage ceiling. This is the formal content of the quality–separation tradeoff (§Discussion) — a choice of query-selection policy purchases a quantitative, structurally-derived leakage ceiling.

Propagating to the wrapper's belief state: the update $M_{W,m+1}$ is a deterministic function of $(M_{W,m}, o_{W,m+1}, q_M, A(q_M))$, so along $G^{\text{op}} \rightarrow q_M \rightarrow (q_M, A(q_M)) \rightarrow M_{W,m+1}$ the data-processing inequality gives

$$I(M_{W,m+1}; G^{\text{op}} | M_{W,m}, o_{W,m+1}) \leq I(A(q_M); G^{\text{op}} | M_{W,m}, o_{W,m+1}) \leq I(q_M; G^{\text{op}}). \quad (13.33)$$

The conditioning set is $(M_{W,m}, o_{W,m+1})$ — deliberately *not* q_M . Conditioning on q_M would (correctly, for the closed processing path) zero the quantity out and miss the selection leak; the unconditional $\kappa_{W_1}^{\text{sel}} = I(A(q_M); G^{\text{op}}) \leq I(q_M; G^{\text{op}})$ is the analyst-independent statement, with q_M drawn under the wrapper's actual operating policy.

Wrapping regime hierarchy. The construction supports three regimes, distinguished by where structural separation lives:

Table 13.4

<i>Regime</i>	<i>Construction</i>	<i>Leakage bound</i>	<i>Leakage source</i>
W₀ (no wrapping)	Raw Class 2 (Partial) or Class 3 (Coupled) component	κ_{W_0} at the component's maximum goal-conditioning sensitivity	No constraint
W₂ (partial wrapping)	One goal-conditioned call per macro-step; structurally typed parsed response routes updates to M_W vs. G_W slots	κ_{W_2} bounded <i>behaviorally</i> — by the component's compliance with the prompted instruction-to-separate; no structural bound	Component's instruction-following fidelity
W₁ (strict wrapping)	Theorem 1 / 2 — separate q_M and q_G calls per macro-step	$\kappa_{W_1}^{\text{sel}} = I(A(q_M); G^{\text{op}}) \leq I(q_M; G^{\text{op}})$ — bounded <i>structurally</i> by the goal-content of the wrapper's query-selection policy (under (C1)+(C2'))	Query-selection policy's goal-content (the goal G^{op} inferable from query content)

W_1 admits a structural bound — the selection-channel quantity $\kappa_{W_1}^{\text{sel}} = I(A(q_M); G^{\text{op}}) \leq I(q_M; G^{\text{op}})$ under (C1)+(C2'); W_2 admits only a behavioral bound from the component's compliance fidelity. The two are different in kind — the structural bound is a property of the wrapper's query-selection policy (the goal-content of the queries it chooses) and is therefore design-controllable; the behavioral bound depends on the component's training and prompt-following. The W_1 structural bound is available exactly when the component carries no goal-correlated cross-call state ((C2')); when

(C2') fails the leak routes through an unobservable channel and W_1 too degrades to a behavioral bound — the boundary derived in `#disc-w1-structural-bound-boundary`.

The $W_0 / W_2 / W_1$ distinction refines the Class 1 (Separated) cell of `#der-directed-separation`: within Class 1 (Separated), **Class-1-by-structure** (natively goal-blind components, or W_1 wrapping) has a structurally derivable directed-separation guarantee; **Class-1-by-behavior** (W_2 wrapping) has only an empirically estimable guarantee that depends on the component's instruction-following.

Epistemic Status.

Conditional on (C1), (C2), and (C3) (or its weakening to a leakage bound). The proofs are short conditional-independence reasoning (Theorem 1) and a single application of the data-processing inequality (Theorem 2); both are standard.

The W_1 structural leakage bound $\kappa_{W_1}^{\text{sel}} = I(A(q_M); G^{\text{op}}) \leq I(q_M; G^{\text{op}})$ is *conditional* on (C1) and (C2'). The data-processing step is standard; the load-bearing condition is (C2') — the Markov chain $G^{\text{op}} \rightarrow q_M \rightarrow A(q_M)$ holds only when the component carries no goal-correlated state across the call boundary. The bound is an upper bound; its gap below $I(q_M; G^{\text{op}})$ is the “accuracy effect” of §Discussion “Two senses of component competence” and is not characterized here (Working Notes).

Max attainable: derived under stated conditions. (C3)'s exact form is a structural ideal that pretrained components (notably LLMs with goal-rich training data) generally satisfy only approximately; the realistic regime is Theorem 2 with κ characterized empirically. (C2)'s exact form is likewise an ideal — a frozen-weights LLM with a conversation cache violates it without any weight adaptation — and where it fails only the behavioral bound of `#disc-w1-structural-bound-boundary` survives.

The wrapping regime hierarchy ($W_0/W_2/W_1$) is a *formulation* — the partition is made by the structural choice of where to place the separation commitment. The leakage bounds within each regime are derived once the regime is fixed.

Discussion.

Quality–separation tradeoff inside Class B. For Class-B components (admitting both goal-blind and goal-conditioned modes), the wrapper has a design choice: how aggressively to restrict q_M to goal-blind content, vs. how much context to allow that may carry goal-correlated information. Maximally goal-blind queries (only the current observation, no context, no history) drive the selection-channel mutual information $I(q_M; G^{\text{op}})$ — and with it the leakage ceiling $\kappa_{W_1}^{\text{sel}}$ — toward zero, but may produce information-poor responses that hurt f_M 's update quality. Maximally informed queries (full history, retrieved context) produce richer responses but raise $I(q_M; G^{\text{op}})$ and therefore the upper bound on $\kappa_{W_1}^{\text{sel}}$. The tradeoff is real and resolved per application — and because $I(q_M; G^{\text{op}})$ is a property of the query-selection policy the wrapper controls, the tradeoff is a design dial, not a property of the component.

Component-admissibility partition. Class A components (goal-blind by design) satisfy (C1) trivially and don't need wrapping in the substantive sense — wrapping for Class A is organizational rather than structural. Class B components (LLMs, hybrid RL with separable value/policy, multi-modal models) are the substantive wrapping case — the wrapper *chooses* to use the goal-blind mode. Class C components (pure end-to-end goal-conditioned policy networks) fail (C1) and are scope-out for the basic theorem. Salvage paths for Class C — null-goal queries, goal-uniform averaging, auxiliary distilled goal-blind heads — exist but cost something (information loss, computation, training).

Resolution of the LLM scope question. The “Class 3 (Coupled) exit” framing — *directed separation violated by goal-conditioned agents (LLMs); handled as architectural scope, not approximation* — is refined by this segment from a scope exit to a constructive route through. Class 3 (Coupled) LLMs are scope-in *for the wrapper construction* (under Class-B admissibility). The cost is paid in residual leakage rate $\kappa_{W_1}^{\text{sel}}$ bounded by the goal-content of the wrapper’s query-selection policy, $I(q_M; G^{\text{op}})$; the tempo cost is established separately in #der-class-coercion-in-composition. Whether this construction yields an operationally useful agent depends on how low the query-selection policy can drive $I(q_M; G^{\text{op}})$ while still licensing useful belief updates — and, where the component carries goal-correlated cross-call state, on whether (C2’) can be enforced at all (#disc-w1-structural-bound-boundary).

Relationship to #hyp-directed-separation-under-composition. The hypothesis is descriptive — when does directed separation hold under composition? This segment provides the constructive answer for the wrapper-around-component special case: directed separation holds whenever the wrapper’s type signatures are respected and (C1)–(C3) hold (or their weakenings). The general N-agent composition question remains a hypothesis; the wrapper-around-component case is now derived.

Wrapping as a truthification mechanism. The wrapping construction is the *rigorous formal version* of what #disc-adversarial-coupling-pressure §”Defensive scaffolding as composition” gestures at informally — peer review, prediction registers, double-entry bookkeeping, adversarial procedure, structured red-teaming. Those external scaffolds are operational mechanisms for *increasing* the modularity of a composite agent in the face of forces that would couple it; the wrapping construction is the structural version of the same operation applied internally rather than externally. Both share the discipline: a goal-blind belief-update query path is structurally enforced (W_1 strict) or behaviorally bounded (W_2), at a definite cost (extra component calls per macro-step plus residual leakage rate). The $W_0/W_2/W_1$ regime hierarchy is the *graded* characterization of how thoroughly the truthification has been applied — W_0 is the un-truthified base state, W_2 behavioral truthification, W_1 structural truthification. Cross-reference: #disc-modularity-state-dynamics is the meta-segment in which the truthification operation sits as one of three operations on the modularity state — alongside strategic self-coupling (#disc-strategic-self-coupling, self-driven-decreasing) and adversarial coupling pressure (#disc-adversarial-coupling-pressure, externally-driven-decreasing). This segment is the canonical *formal* instance of the truthification operation at the component level ($W_0 / W_1 / W_2$ regime hierarchy), paired with the composite-level *defensive-scaffolding* instance from #disc-adversarial-coupling-pressure §”Defensive scaffolding as composition”. The connection is also surfaced in #impl-composition-machinery §”Class-coercion as truthification mechanism.”

Two senses of component competence: world-simulation vs goal-extraction. A subtle decomposition of *component competence* matters for how (C3) and the selection-channel bound $\kappa_{W_1}^{\text{sel}} = I(A(q_M); G^{\text{op}}) \leq I(q_M; G^{\text{op}})$ scale with the component’s capability. A high-capability component has at least two distinct competences that the classification machinery treats asymmetrically:

1. **World-simulation fidelity** — accuracy of A ’s response to the explicit content of a query, including counterfactual content. Sharpens f_M ’s update quality once a query is in hand.
2. **Input-structure extraction** — competence at inferring latent structure *from the query’s surface form*, including inferring the operator’s goal G^{op} from query patterns. This is the channel (C3) names.

The architectural classification of #der-directed-separation is *blind to (1) and defined by (2)*: class membership turns on whether the goal can causally reach f_M , not on how accurately f_M then computes given goal-blind input. For pretrained components the two competences correlate empirically — the same training that produces world-simulation fidelity also produces input-structure extraction — but conceptually they are distinct, and the distinction is load-bearing for (C3).

The two competences act on the two sides of the selection-channel bound. The bound’s *ceiling* is $I(q_M; G^{\text{op}})$ — the goal-content present in the query, fixed by the wrapper’s selection policy; the *realized* leak is $I(A(q_M); G^{\text{op}})$ — how much of that the response actually carries. The gap between them is governed by competence:

- *Accuracy effect (widens the gap, lowers the realized leak).* A more competent world-simulator returns posteriors closer to the true conditional $P(\cdot \mid q_M)$, which is goal-independent by definition; it leaves the goal-content sitting in q_M unexploited in its response, so $I(A(q_M); G^{\text{op}})$ can sit well below the ceiling $I(q_M; G^{\text{op}})$.
- *Inference effect (closes the gap, raises the realized leak).* A more competent input-structure extractor uses any query-content / goal-content correlation to infer G^{op} from the goal-blind query and bend its response accordingly, driving $I(A(q_M); G^{\text{op}})$ up toward the ceiling — even when no goal argument is structurally exposed.

For pretrained components where the two competences scale together (the typical case for LLMs trained on goal-rich corpora), where the realized leak sits between 0 and the ceiling in the fidelity limit is *empirical and component-specific*, not structural. The selection-channel ceiling is design-controllable (drive $I(q_M; G^{\text{op}})$ down by goal-blind query selection); the position within that ceiling is the accuracy-vs-inference balance, and separating the two appears to open more science than it closes.

Several lines worth pulling, kept open here rather than queued elsewhere so future cycles do not lose them:

- *Decorrelation by construction.* If world-simulation and input-structure extraction can be structurally decoupled at training time — by architectural separation of *what does the world look like given q* from *what is the latent context that produced q* — a Class 3 (Coupled) component could satisfy (C3) even at high world-simulation fidelity. The classification would then refine into “Class 3 but C3-respecting” vs “Class 3 and C3-violating” sub-cases that the current scalar κ_{W_1} does not surface, and the C3-respecting branch may be a more interesting design target than wrapping a C3-violating component.
- *Counterfactual capability is bi-side.* Counterfactual simulation (Pearl Level 2) requires both competences — accurate world-simulation under interventional queries *and* extraction of the intervention’s intended scope from the query. Whether high counterfactual capability tends to raise or lower realized κ in practice is an open empirical question; the answer depends on how the two competences decompose in the specific component.
- *Decomposed empirical estimator.* The existing $\hat{\kappa}_{\text{processing}}$ estimator from #der-directed-separation probes the *net* response divergence under varied goal-priming. An estimator that holds explicit query content fixed while varying only the implicit operator context might isolate the inference-effect contribution, giving a sharper diagnostic than the aggregate κ and surfacing whether a given high-fidelity component sits near $\kappa \rightarrow 0$ (accuracy-dominated) or near κ ceiling (inference-dominated).
- *M4 amplification specificity.* The amplification of #disc-modularity-state-dynamics operations (strategic self-coupling, adversarial coupling pressure) under high component fidelity may track competence (2) specifically rather than competence (1). If so, M4 attack-surface measurement could be sharpened by measuring the inference effect rather than aggregate capability — a more competent world-simulator is not, *by that fact alone*, a more potent vector for the coupling-increasing operations.

The decomposition is named here at *discussion-grade* recognition — kept open as a place to pull on rather than as a closed result. Whether it lifts to *robust qualitative* or higher depends on whether the two competences can be operationalized as separately measurable quantities and whether the architectural-decorrelation conjecture holds for any realizable training regime.

Findings.

Constructive Directed Separation via Wrapping. **Brief:** When you have a component (like an LLM) whose belief-update and goal-conditioning are entangled in a single forward pass, you can build a scaffold around it that maintains explicit, separate stores for what the system believes and what it wants. The structural rule is that belief updates only see queries to the component that don't include the goal as input. Under reasonable conditions on the component, the wrapped system is goal-blind in its belief updates *by construction* — even though the underlying component isn't. The cost shows up as a residual leakage from the component's pretraining (the component might still infer the goal from query content, even when the goal isn't explicit in the input). Two practical regimes appear: strict wrapping with separate goal-blind and goal-conditioned calls (theoretically clean, with a structural leakage bound), and partial wrapping with one goal-conditioned call whose response is parsed into separate update fields (operationally common, with only a behavioral leakage bound — depending on the component's instruction-following fidelity rather than its query structure).

Impact: Promotes #hyp-directed-separation-under-composition to derived (in the wrapper-around-component special case). Refines the Class 1 (Separated) cell of #der-directed-separation with a structural-vs-behavioral sub-distinction (W_1 vs. W_2). Resolves the LLM scope question — Class 3 (Coupled) components are scope-in for the wrapper construction at a measurable cost, not scope-out. The composition-level consequences (wrapper as valid AAT composite agent, persistence-template inheritance, tempo cost) are derived in the companion segment #der-class-coercion-in-composition.

Novelty Claim: *Claim integration* of POMDP / cognitive-architecture prior art with the AAT Class 1/2/3 (Separated/Partial/Coupled) directed-separation taxonomy, plus the $W_0/W_2/W_1$ regime hierarchy that surfaces the structural-vs-behavioral leakage distinction and the LLM-specific (C1)–(C3) admissibility/leakage conditions. The wrapping move itself is rediscovery of patterns established in POMDP theory (Bayesian belief-update is goal-blind by construction) and cognitive architectures (modular agent design with separated belief/goal/action state, four decades). AAT's contribution is the structural-leakage analysis at the directed-separation level and the regime hierarchy that names where the separation guarantee lives.

Related Work:

Table 13.5

<i>ASF concern</i>	<i>Prior-art language</i>	<i>Relationship / Positioning</i>
Goal-blind belief-update by construction	Astrom 1965, "Optimal control of Markov processes with incomplete state information," <i>J. Math. Anal. Appl.</i> 10; Kaelbling, Littman, Cassandra 1998, "Planning and acting in partially observable stochastic domains," <i>Artificial Intelligence</i> 101	<i>formal antecedent</i> — POMDP belief-state filters are goal-blind by construction; the wrapping move recapitulates this in the AAT vocabulary. The closest formal prior art for the directed-separation guarantee.
Modular agent design with separated belief/goal/action	Newell 1990, <i>Unified Theories of Cognition</i> ; Laird 2012, <i>The Soar Cognitive Architecture</i> ; Anderson 2007, <i>How Can the Human Mind Occur in the Physical Universe?</i> ; Sun 2016 <i>Anatomy of the Mind</i> (CLARION); Baars 1988 <i>A Cognitive Theory of Consciousness</i> / Dehaene 2014 <i>Consciousness and the Brain</i> (Global Workspace)	<i>formal antecedent</i> — cognitive architectures have done modular agent design with separated belief/goal/action state for 40+ years. The W_1 wrapping move is essentially the per-cycle commitment that cognitive architectures make at the system level.

<i>ASF concern</i>	<i>Prior-art language</i>	<i>Relationship / Positioning</i>
Tool-using language-model agent frameworks	Yao et al. 2022, “ReAct: Synergizing reasoning and acting in language models”; Shinn et al. 2023, “Reflexion: language agents with verbal reinforcement learning”; Park et al. 2023, “Generative Agents: Interactive Simulacra of Human Behavior”; Packer et al. 2023, “MemGPT: Towards LLMs as operating systems”; Schick et al. 2023, “Toolformer: language models can teach themselves to use tools”	<i>empirical instantiation</i> — practical wrappers around language-model substrates. Most fall in W_2 (partial wrapping / output-structuring); Generative Agents’ observation→memory step is the closest empirical instance of W_1 . AAT’s regime hierarchy gives these constructions a structural reading.
Hybrid deliberative/reactive architectures (prior art for class-coercion)	Gat 1992, <i>Integrating planning and reacting in a heterogeneous asynchronous architecture for controlling real-world mobile robots</i> (AAAI) ⁴ ; Simmons 1994, <i>Structured control for autonomous robots</i> (IEEE Trans. Robotics 10:34) ⁵ ; Au 2004, planner wrappers with external-query management ⁶	<i>formal antecedent</i> — hybrid deliberative/reactive architectures orchestrate reactive (entangled) layers beneath deliberative (separated) planners, structurally precedent for the wrapping construction’s “scaffold an entangled component to recover separated behaviour at the wrapper level” move. AAT’s contribution is <i>not</i> the wrapping move itself but the <i>theorem-shaped</i> wrapper-level directed-separation guarantee (Theorem 1 + Theorem 2) plus the explicit selection-channel leakage bound $\kappa_{W_1}^{se} = I(A(q_M); G^{op}) \leq I(q_M; G^{op})$ and the Brooks’s-Law tempo cost in #der-class-coercion-in-composition

Search Log:

- 2026-05-09 (*targeted*): Web + training-data search across POMDP / cognitive-architecture / scaffolded-LLM threads. Verdict: **substantial overlap** with the POMDP and cognitive-architecture lines as the closest formal prior art. AAT’s contribution is the structural-leakage analysis and regime hierarchy rather than novelty in the wrapping move itself.
- 2026-05-09 (*intuition-only*, prior to the targeted search): adjacent literatures expected to host prior art were active inference (Markov blankets), control theory (approximate simulation), and scaffolded-LLM frameworks. The targeted search confirmed all three and added the POMDP and cognitive-architecture lines as the formal precedents.

Working Notes.

- **Empirical κ measurement.** The selection-channel ceiling $I(q_M; G^{op})$ is estimable by varying the latent operator goal across a fixed query-selection policy and measuring the divergence of the resulting query distribution; the realized leak $I(A(q_M); G^{op})$ requires sampling responses under varied G^{op} at fixed selection policy. Both need an operationalization of the latent G^{op} —

⁴Citation surfaced 2026-05-22 from the Track E catalog at ref/prior-art-analysis/ (intermediate work artifacts that captured Pillar-style prior-art searches). Catalog has more verification support than raw Undermind synthesis but less than full primary-source reading. Verification queued with the BG2 cluster.

⁵Citation surfaced 2026-05-22 from the Track E catalog at ref/prior-art-analysis/ (intermediate work artifacts that captured Pillar-style prior-art searches). Catalog has more verification support than raw Undermind synthesis but less than full primary-source reading. Verification queued with the BG2 cluster.

⁶Citation surfaced 2026-05-22 from the Track E catalog at ref/prior-art-analysis/ (intermediate work artifacts that captured Pillar-style prior-art searches). Catalog has more verification support than raw Undermind synthesis but less than full primary-source reading. Verification queued with the BG2 cluster.

likely the same construct the $\hat{\kappa}_{\text{processing}}$ estimator of $\#der\text{-directed-separation}$ uses. For a fixed component, the ceiling is a property of the wrapper's q_M -selection policy — narrower (more goal-blind) queries reduce it, richer queries raise it. The empirical instantiation, and a characterization of the gap between realized leak and ceiling (the accuracy effect), is open follow-on; see also the decomposed-estimator spike candidate below.

- **Sub-type-aware wrapping regime selection.** Knowing only that a component is Class 2 (Partial) is insufficient to choose W_1 vs W_2 ; the right regime depends on the un-wrapped component's coupling *form* per the (stage $\times$ source $\times$ form) sub-typology of $\#disc\text{-partial-coupling-pathways}$: content-form sub-types admit W_2 (post-hoc response structuring); process-form sub-types require W_1 (structural goal-blind query path) or are not coercible to Class 1 without pipeline access. Additionally, Σ_t -source coupling (per $\#der\text{-belief-strategy-attractor}$) can undermine W_1 's structural commitment via strategy-context leakage in stateful components — even with no G_W in the query, the component's internal Σ may be influenced by historical query content and then suppress the gain on subsequent goal-blind queries; a Σ -channel-suppressed W_1 (holding strategic context fixed across calls, or stripping Σ -content from queries) may be required. The $W_0/W_1/W_2$ hierarchy is the agent-level analog of the stage-level form distinction in the sub-typology — $W_2 \leftrightarrow \text{content}$; $W_1 \leftrightarrow \text{process-with-pipeline-access}$ — making the structure-vs-behavior refinement at Class 1 the agent-level shadow of a Class 2 axis.
- **Compositional wrapping (wrapper-of-wrapper).** How leakage rates compose under iterated wrapping is open. Conjecture: additive in KL ($\kappa_{\text{outer}} \leq \kappa_{\text{inner}} + \kappa_{\text{outer-shell}}$) by data-processing inequality applied at each level, but tightness is unclear.
- **Behavioral compliance axiom for W_2 .** κ_{W_2} has no structural bound; it depends on the component's instruction-following fidelity. Whether a behavioral-compliance axiom (assuming the component honestly attempts to follow structural-separation instructions) yields a bound is an open hypothesis. If so, it would be hypothesis-grade rather than derived.
- **Identifying the regime in the wild.** Practical scaffolded-LLM frameworks (ReAct, Reflexion, MemGPT, etc.) almost universally implement W_2 . Distinguishing W_2 from W_1 in a deployed system requires inspection of the per-cycle query structure — does f_M 's update path receive a query that contains G_W or not? This is the diagnostic question.
- **Segment split provenance (2026-05-11).** This segment was bifurcated from a combined “class coercion” derivation. Claim A (directed separation at the wrapper level) lives here; Claim B (wrapper as valid AAT composite agent — (A1)–(A4) verification, persistence-template inheritance, Brooks's-Law tempo cost) lives in $\#der\text{-class-coercion-in-composition}$ (which declares this segment as prerequisite). The split reflects FORMAT.md Gate 1 discipline: this segment's depends list ($der\text{-directed-separation}$, $def\text{-agent-environment}$) reflects exactly what the directed-separation theorem actually requires. The composition-level dependencies ($form\text{-composition-closure}$, $deriv\text{-sector-condition}$, $result\text{-sector-persistence-template}$, $der\text{-tempo-composition}$) are Claim B's load and now live with Claim B.
- Reasoning-trail provenance: spike directories at `spikes/class-coercion-wrapping/` and `spikes/temporal-nesting-rg/` carry the working-out of these results.
- **W_1 leakage-bound — regression guard.** Do not regress the selection-channel bound $\kappa_{W_1}^{\text{sel}} = I(A(q_M); G^{\text{op}}) \leq I(q_M; G^{\text{op}})$ back to the form $\kappa_{W_1} \leq I(A(q_M); G_W \mid q_M)$ that the body previously carried: that quantity is circular and vacuous — it bounds κ by an average of κ , and evaluates identically to zero for the stateless oracle the theorem models, because conditioning on q_M closes the only channel it measures. The fix relocated the goal variable from the wrapper register G_W to the latent operator goal G^{op} read off query content and dropped the conditioning (data-processing along $G^{\text{op}} \rightarrow q_M \rightarrow A(q_M)$ under (C1)+(C2')); the embedded no-go — structural bound available iff (C2') holds, otherwise behavioral-only — is $\#disc\text{-w1-structural-bound-boundary}$. Reasoning trail (filed): `spikes/.integrated/spike-w1-leakage-vacuity-2026-05-31.md`, `spikes/.integrated/spike-w1-w2-boundary-intuition-2026-05-31.md` (cycle narrative: CHANGELOG 2026-05-31).
- **Migration note (2026-05-09 GUC rename):** Class 2 $\leftrightarrow$ Class 3 swap. Pre-2026-05-09: Class 2 = fully merged, Class 3 = partially modular. Post: Class 2 = Partial, Class 3 = Coupled. Three independent axes in this segment: (a) GUC Class 1/2/3 — renamed and swapped; (b) $W_0/W_1/W_2$ wrapping regimes — UNCHANGED; (c) Class A/B/C component-admissibility partition — UNCHANGED.
- **Track E surface-back of catalog citations (2026-05-22).** One Related Work entry added 2026-05-22 from Track E catalog at `ref/prior-art-analysis/05-directed-separation.md` (Pillar 4): hybrid deliberative/reactive architectures (Gat 1992, Simmons 1994, Au 2004) as prior art for class-coercion. Marked [[^]cat-2026-05-22] for verification-deferred attribution; primary-source verification queued in the BG2 cluster. The catalog citations had prior Pillar-style search support but were not verified by the current executor at landing time.

- **Spike candidate (PROPOSED Tier 3, 2026-05-28): Decomposed empirical estimator separating accuracy-dominated from inference-dominated κ regimes.** The aggregate $\hat{\kappa}_{\text{processing}}$ estimator of $\#der\text{-directed-separation}$ probes net response divergence under varied goal-priming. A decomposed estimator that holds explicit query content fixed while varying only implicit operator context (system-prompt scaffolding, query metadata, conversational history) might isolate the inference-effect contribution from the accuracy-effect contribution — surfacing whether a given high-fidelity component sits near $\kappa \rightarrow 0$ (accuracy-dominated) or near a κ ceiling (inference-dominated). Substance home: §Discussion “Two senses of component competence: world-simulation vs goal-extraction”. → [PROPOSED Tier 3](#).
- **Spike candidate (PROPOSED Tier 2, 2026-05-28): Decorrelation-by-construction — Class 3 but C3-respecting sub-cases.** Architectural-research question: can world-simulation fidelity and input-structure extraction be structurally decoupled at training time (training-objective decomposition; architectural separation of *what does the world look like given q* from *what is the latent context that produced q* ; data-curation discipline)? If so, a Class 3 (Coupled) component could satisfy (C3) at high world-simulation fidelity, and the architectural classification of $\#der\text{-directed-separation}$ would refine into “Class 3 but C3-respecting” vs “Class 3 and C3-violating” sub-cases that the scalar κ_{W_1} does not currently surface — the C3-respecting branch may be a more interesting design target than wrapping a C3-violating component. Theory-edge / speculative; substance home: §Discussion “Two senses of component competence: world-simulation vs goal-extraction”. → [PROPOSED Tier 2](#).

Incidental audit gold (lift 2026-05-30). Cross-audit “wandering thoughts” / §14-ideation harvested from the de-novo auditors’ working dirs. **Coverage:** only one dir reached a digested reflection on this segment (526815) — the deeply-mathematical critique pass; the other contributing dirs stopped before this segment. **No pedagogical / analogy / figure gold surfaced here** — 526815’s reflection is entirely *certified-track correctness findings* (the F127–F131 cluster on the leakage analysis), which are routed to the findings track, not lifted as Working-Notes gold. The one durable non-finding signal worth preserving for the later Brief/Discussion pass:

Belongs elsewhere / convergence signal.

- **Auditor flagged the $\varepsilon_{\text{track}}$ / $\varepsilon_{\text{coerce}}$ / κ distinction as a contribution to preserve** — “the distinction ... is excellent and should be preserved downstream” (Codex/Claude, AUDIT-WORKING-526815, “Watch”). A signal about which of this segment’s apparatus a fresh reader found most load-bearing; useful when the leakage analysis is eventually given Brief-field plain-language framing.

Off-ramp note: 526815 also raised a substantive *structural* challenge to the W_1 leakage bound — that $\kappa_{W_1} \leq I(A(q_M); G_W | q_M)$ is identically zero for a stateless component under exact conditioning on q_M , so the bound as stated may not measure the leakage it intends (the real channels being query-content correlation, hidden component state, conversation history). This was a strengthen-first / correctness matter for the chapter’s load-bearing construction, **not** Working-Notes gold — flagged in the lift report for the findings/adjudication track. **Resolved** via the strengthen-first spike: the bound was replaced by the selection-channel quantity $\kappa_{W_1}^{\text{sel}} = I(A(q_M); G^{\text{op}}) \leq I(q_M; G^{\text{op}})$ and the embedded no-go landed as $\#disc\text{-}w1\text{-structural-bound-boundary}$ (see the W_1 leakage-bound correction provenance note above).

13.5 Derived: Class Coercion in Composition

Under the wrapper construction of $\#der\text{-class-coercion-via-wrapping}$, the wrapped system W is a valid AAT composite agent: it satisfies (A1)–(A4) of $\#form\text{-composition-closure}$, inherits the sector-persistence template from $\#result\text{-sector-persistence-template}$, and incurs a tempo cost in the Brooks’s-Law form of $\#der\text{-tempo-composition}$. This segment establishes those composition-level consequences; the directed-separation guarantee that motivates the construction is established in the prerequisite segment.

Three additional design constraints on the wrapper itself are required for (A2)–(A4) to hold at the wrapper level: that the wrapper commits to a prediction map so macro-mismatch is well-defined; that its belief-update map supports a gain interpretation per $\#def\text{-adaptive-tempo}$; and that the belief-update map satisfies the sector condition with positive correction rate. All three are automatic for Tier-1 belief-update maps via $\#deriv\text{-sector-condition}$ and $\#der\text{-gain-sector-}$

bridge; Tier-2/3 cases inherit the corresponding tier-restricted scope. (A1) holds by construction — the wrapper *builds* in the AAT-shaped state $X_W = (M_W, G_W)$.

Persistence at the wrapper level requires the wrapper’s correction rate to exceed total effective disturbance, which now has *two* contributions: external environmental disturbance acting through the wrapper’s observation channel, and internal disturbance from the component’s response variance, bounded by the variance of the component’s responses to goal-blind queries. The wrapping construction is honest about what it does *not* buy: the component’s internal non-determinism still acts as a disturbance on the wrapper level. What it eliminates is the *belief-goal coupling* that the directed-separation classification identifies as the load-bearing failure mode for Class 3 (Coupled) substrates.

The tempo cost is a **Brooks’s-Law instance**: the wrapper makes $K \geq 2$ component calls per macro-step (one goal-blind for model update, one goal-conditioned for purposeful-state update; more in richer wrapper designs), so the wrapper’s macro-event rate is at most $1/K$ of the component’s underlying rate. The Brooks’s-Law inequality of #der-tempo-composition predicts exactly this — more component calls per macro-step means lower macro-tempo, with the coordination overhead specific to maintaining the wrapper’s (M_W, G_W) state separately from the component’s internal state. The structural reading: a wrapped system gets directed separation (Class 1 status, the full Part II machinery applies) at the *structural price of admission* of (at minimum) halved macro-tempo plus an irreducible internal-disturbance term. This is the framework’s quantitative content for the cost of converting a Class 3 (Coupled) substrate into a Class 1 (Separated) composite.

Formal Expression.

Setup (inherited). The wrapper structure, the four type-signed update components (q_M, q_G, f_M, f_G) , the macro-step schedule ($K \geq 2$ component calls per macro-step), and the admissibility conditions (C1)–(C3) on the component A are defined in #der-class-coercion-via-wrapping. This segment uses that setup without redefinition.

Wrapper-design constraints. For (A2)–(A4) of #form-composition-closure to hold at the wrapper level, the wrapper’s update structure must satisfy three additional constraints:

(D-A2) The wrapper commits to a prediction map $\hat{o}_W : \mathcal{X}_M \times \mathcal{A}_W \rightarrow \mathcal{O}_W$ so that macro-mismatch $\delta_W = o_W - \hat{o}_W$ is well-defined.

Conditions (wrapper-design).

(D-A3) f_M supports a gain interpretation per #def-adaptive-tempo. Holds for Tier-1 belief-update maps — Bayesian on exponential families, gradient on strongly convex losses, linear-PD with bounded gain. Tier-2/3 cases inherit the corresponding tier-restricted scope from #deriv-sector-condition.

(D-A4) f_M satisfies the sector condition with positive correction rate. Automatic for Tier-1 belief-update maps via #deriv-sector-condition Prop A.1 and #der-gain-sector-bridge.

(A1) of #form-composition-closure holds by construction — $X_W = (M_W, G_W)$ has the AAT form because the wrapper *builds it in*.

Theorem: Wrapper as valid AAT composite agent.

Under (D-A2)–(D-A4) and the directed-separation conditions of #der-class-coercion-via-wrapping, the wrapper W satisfies (A1)–(A4) of #form-composition-closure and therefore qualifies as an AAT composite agent.

Derived (wrapper-as-composite, from (D-A2)+(D-A3)+(D-A4)).

Proof. (A1) by construction of the type signatures; the wrapper’s state is $X_W = (M_W, G_W)$ in the AAT shape. (A2) under (D-A2): the prediction map closes the mismatch definition. (A3) under (D-A3): Tier-1 belief-update maps inherit gain interpretation via #deriv-sector-condition Prop A.1, with the Tier-2/3 lifts deferring to the tier-restricted scope

from `#form-composition-closure`'s bridge-lemma classification. (A4) under (D-A4): the sector condition transfers from the (Tier-1) belief-update map to the wrapper level via `#der-gain-sector-bridge`. ■

Inheritance of the persistence template. Under (D-A4), the wrapper inherits `#result-sector-persistence-template` at the wrapper level: persistence holds when $\alpha_W R_W > \rho_W$. The wrapper-level effective disturbance has two contributions: external environmental disturbance ρ_{ext} acting through o_W , and internal disturbance from the component's response variance, ρ_{int} , bounded by the variance of A 's responses to goal-blind queries. Total: $\rho_W = \rho_{\text{ext}} + \rho_{\text{int}}$. Persistence at the wrapper level requires $\alpha_W R_W > \rho_{\text{ext}} + \rho_{\text{int}}$.

Tempo cost — Brooks's-Law instance. The wrapper makes $K \geq 2$ component calls per macro-step (more in richer wrapper designs). If the component's nominal call rate is ν_A , the wrapper-level macro-update rate is $\nu_W = \nu_A/K$. By `#der-tempo-composition`,

$$\mathcal{T}_W \leq \mathcal{T}_A^{\text{nominal}} - C_{\text{coord}}^{\text{wrap}} \quad (13.34)$$

where $C_{\text{coord}}^{\text{wrap}}$ is the coordination overhead specific to the wrapping construction — the tempo consumed by maintaining the wrapper's (M_W, G_W) state separately from the component's internal state. This is the cost of class coercion paid in tempo: the same Brooks's-Law form whose general statement is in `#der-tempo-composition`. Adding state-management infrastructure reduces realized external tempo even when the underlying component's compute rate is unchanged.

Epistemic Status.

Conditional on the wrapper-design constraints (D-A2), (D-A3), (D-A4), and on the prerequisite directed-separation result of `#der-class-coercion-via-wrapping`. The proof is direct verification using inheritance from `#deriv-sector-condition` and `#def-adaptive-tempo`; both are standard. (D-A4) is automatic for Tier-1 belief-update maps; for Tier-2/3 belief-update maps the wrapper inherits the tier-restricted scope from `#deriv-sector-condition` and `#form-composition-closure`'s bridge-lemma classification.

Max attainable: derived under stated conditions. The composition-level claim is parametrized by the belief-update map's tier; Tier-1 cases inherit cleanly, Tier-2/3 cases inherit under additional restrictions.

Discussion.

Coercion-distance vs. tracking-distance. The closure-defect ε^* in `#form-composition-closure` is read as a *fidelity* measure: how well the macro-system tracks the projected micro-system. In the wrapping construction, the wrapper deliberately changes the underlying component's behavior to enforce separation — the wrapper *should* differ from the unwrapped component in the goal-conditioning direction. Two distinct quantities are involved:

- $\varepsilon_{\text{track}}^*$ — the standard fidelity quantity from `#form-composition-closure`. Used in the bridge lemma to bound trajectory error.
- $\varepsilon_{\text{coerce}}^*$ — the wrapping-specific quantity measuring behavioral divergence between the wrapped system and the unwrapped component. Higher means more aggressive coercion.

For wrapper analyses: the persistence-template inheritance (above) uses $\varepsilon_{\text{track}}^*$. Cost-of-class-coercion analyses use $\varepsilon_{\text{coerce}}^*$. The leakage analysis from `#der-class-coercion-via-wrapping` uses κ — a KL-divergence on response distributions, distinct from both trajectory-error quantities. Conflating them propagates downstream confusion.

Form-preservation reading. Read in form-preservation language, (A1)–(A4) of #form-composition-closure are the requirement that AAT’s form survive coarse-graining: macro must itself be AAT. The wrapping construction is the *constructive* version — when the underlying component doesn’t preserve form on its own (Class 2 (Partial) or Class 3 (Coupled)), the wrapper enforces form-preservation at the macro level by structural commitment of the type signatures (from #der-class-coercion-via-wrapping) plus the wrapper-design constraints here. The closure-defect $\varepsilon_{\text{track}}^*$ then plays the role of distance from the form-preserved fixed point in the constructive setup.

Findings.

Wrapper as Valid AAT Composite Agent with Brooks’s-Law Tempo Cost. **Brief:** Once you’ve wrapped a coupled component to get directed separation (the structural rule from #der-class-coercion-via-wrapping), is the wrapped system actually a valid AAT agent that AAT’s machinery applies to? The answer is yes, under reasonable design constraints on the wrapper’s prediction map, gain interpretation, and sector behavior. This means the wrapped system inherits the persistence-condition machinery — the wrapper’s own version of “can the agent keep up with disturbance” applies, with disturbance now decomposing into the external environment’s drift plus the underlying component’s response variance. There’s a cost, and it’s Brooks’s Law in tempo coordinates: every additional component call per cycle (needed to maintain separate belief and goal stores) divides the macro-update rate. The wrapped system runs slower than the component would running alone, by a structurally determined factor.

Impact: Closes the loop on whether the wrapping construction yields a *usable* AAT agent: yes, the wrapper inherits #result-sector-persistence-template and #der-tempo-composition at the wrapper level. Makes the tempo cost of class coercion explicit and quantifiable — it’s a Brooks’s-Law instance, the same general form that governs all AAT compositions, not a special-case tax. Provides the basis for analyzing logogenic-substrate wrappers (e.g., PROPRIUM’s auxilia hierarchy, agentic-tft’s CONTEXTUALIZE-CHOOSE phase separation) as valid composite agents with computable persistence and tempo characteristics.

Novelty Claim: *Claim integration* of the AAT sector-Lyapunov persistence template, Brooks’s-Law tempo accounting, and the form-composition-closure (A1)–(A4) discipline, applied to the wrapper-around-component construction. The composition-level inheritance is straightforward given the persistence-template and tempo-composition machinery; the contribution is establishing that the wrapped system is in the right shape for those results to apply, and naming the Brooks’s-Law form of the coercion cost.

Related Work:

Table 13.6

<i>ASF concern</i>	<i>Prior-art language</i>	<i>Relationship / Positioning</i>
Predictive-loss bounds under coarse-graining of MDPs	Ravindran-Barto 2004, “An algebraic approach to abstraction in reinforcement learning”; Taylor, Precup, Panangaden 2008, “Bounding performance loss in approximate MDP homomorphisms,” NeurIPS; Abel, Hershkowitz, Littman 2016, “Near optimal behavior via approximate state abstraction,” ICML; Subramanian, Sinha, Seraj, Mahajan 2020, “Approximate information state for approximate planning,” arXiv:2010.08843; Congeduti, Mey, Oliehoek 2020, “Loss bounds for approximate influence-based abstraction,” arXiv:2011.01788	<i>adjacent literature</i> — control-theoretic predictive-loss bounds. AAT connects to this cluster via #form-composition-closure’s bridge lemma at the wrapper level.

<i>ASF concern</i>	<i>Prior-art language</i>	<i>Relationship / Positioning</i>
Compositional structure for nested agents	Smithe 2024, “Structured Active Inference,” arXiv:2406.07577; Capucci, Gavranović, Hedges, Rischel 2022, “Towards foundations of categorical cybernetics,” ACT 2021	<i>adjacent literature</i> — categorical / lens framing of compositional agents. The wrapper construction is consistent with the lens framing; an alternative reading in categorical-systems-theory terms is available.
Form-preservation under coarse-graining	Friston 2019 <i>J. Theor. Biol.</i> , “On Markov blankets and hierarchical self-organisation”; Friston, Heins, Verbelen, Da Costa et al. 2025 <i>Front. Network Physiology</i> , “From pixels to planning: scale-free active inference”	<i>adjacent literature</i> — Friston’s framing of the renormalization group as form-conservation under coarse-graining. The wrapping construction reads as a constructive form-preservation move.
IB-Lagrangian semigroup composition	Mehta, Schwab 2014, “An exact mapping between the Variational Renormalization Group and Deep Learning,” arXiv:1410.3831; Kline, Palmer 2022, “Gaussian Information Bottleneck and the Non-Perturbative Renormalization Group,” PMC8967309	<i>adjacent literature</i> — IB-as-RG. The wrapper’s (P1) information-preservation condition (from #form-composition-closure) is IB-shaped; the semigroup composition rule is the closest analog of “AAT form preserved under iterated wrapping.”
Singular-perturbation as unified RG	Chen, Goldenfeld, Oono 1996, “Renormalization group and singular perturbations,” <i>Phys. Rev. E</i> 54:376; Chiba 2009, <i>SIAM J. Appl. Dyn. Syst.</i> 8:1066	<i>formal antecedent</i> — RG framework subsumes singular perturbation theory. The $K_c \gg 1$ timescale-separation regime in #form-composition-closure invokes this for free.

Search Log:

- 2026-05-09 (*targeted*): Search across MDP-homomorphism / hierarchical-control / categorical-systems-theory / FEP-RG / IB-RG / singular-perturbation-RG threads. Verdict: **substantial overlap** with the MDP-homomorphism and categorical-cybernetics lines as the closest formal prior art for composition under coarse-graining. AAT’s contribution is the specific Brooks’s-Law tempo form and the persistence-template inheritance pattern.

Working Notes.

- **Detailed tempo accounting for canonical wrapper architectures.** $C_{\text{coord}}^{\text{wrap}}$ for ReAct-shape, Reflexion-shape, PROPRIUM-shape wrappers is an empirically-relevant computation deferred from this segment’s promotion. The general bound (Brooks’s-Law form via #der-tempo-composition) is established here; specific architectural breakdowns are follow-on.
- **Connection to ELI-specific structure in 04-eli-core/.** Most ELI-specific content (sovereignty axes, accountability infrastructure, identity factors, substrate-independence) is *added structure* beyond what the composition-level inheritance provides — this segment establishes the wrapper as a valid composite agent; ELI work is what happens on top of that substrate.
- **Segment split provenance (2026-05-11).** This segment was bifurcated from a combined “class coercion” derivation. The prerequisite segment #der-class-coercion-via-wrapping carries the directed-separation claim (Claim A: wrapper is Class 1 (Separated) by construction); this segment carries the composition-level claim (Claim B: wrapper is a valid AAT composite agent with persistence-template inheritance and Brooks’s-Law tempo cost). The split reflects FORMAT.md Gate 1 discipline: each segment’s depends list reflects exactly what its claim actually requires.

Incidental audit gold (lift 2026-05-30). Cross-audit “wandering thoughts” / §14-ideation harvested from the de-novo auditors’ working dirs. **Coverage:** only one dir reached a digested reflection on this segment (526815) — the deeply-mathematical critique pass; other contributing dirs stopped before it. **No pedagogical / analogy / figure gold surfaced** — 526815’s reflection is entirely *certified-track correctness findings* (the F132–F138 cluster), routed to the findings track, not lifted as Working-Notes gold. The non-finding signal worth preserving:

Belongs elsewhere / convergence signal.

- **The Claim-A / Claim-B split was read as a clean, correct architectural move** — “the intended split is good: directed separation lives in the prior segment; composition-level inheritance lives here” (Codex/Claude, AUDIT-WORKING-526815). Affirms the 2026-05-11 bifurcation recorded above.

Off-ramp note: 526815 raised a substantive *gating* challenge — that satisfying (A1)–(A4) is *not* the full #form-composition-closure criterion (which also requires #scope-composite-agent satisfaction and projection admissibility (P1)–(P3)), so the “valid AAT composite agent” conclusion needs the missing scope and projection gates plus a full-loop sector argument rather than an f_M -only one (F132–F134). This is a strengthen-first / scope-honesty matter for a load-bearing claim, **not** gold — flagged in the lift report for the findings/adjudication track.

13.6 Discussion: Additional Implications & Discussion

The chapter is where AAT’s per-agent machinery lifts to the composite level. Closure defect supplies the criterion for when a group of agents legitimately counts as one larger agent; the bridge lemma converts predictive defect into trajectory error; the wrapping construction supplies the constructive route from Class 3 (Coupled) components to Class 1 (Separated) composites; and the closed-form critical-mass derivation unifies four previously-separate persistence inequalities as signed special cases of a single composite-sector condition. The chapter’s two segments with ## Findings carry the load-bearing technical content well; the implications-side payoff is the cross-segment composition — how the four pieces produce composite-level results that no single segment carries — and what those results buy for 03-11m-core/, AI-safety architecture, and the framework’s identifiability-floor pattern at the composition layer.

Closure defect, bridge lemma, and the meaning of composition.

#form-composition-closure’s Findings cover the closure-defect / bridge-lemma machinery in full: when does a group of interacting agents legitimately count as one larger agent? — answered by an inequality, with explicit failure modes on both sides. The closure-defect ε^* is the minimum residual prediction error from collapsing the micro-system into an AAT-shaped macro-description; the bridge lemma converts that prediction-level error into a trajectory-level guarantee bounded by $\varepsilon^* \nu_c / \alpha_c$ under sector-bounded macro-correction plus strong monotonicity of the macro-update. The implication worth surfacing beyond the segment’s own statement is *what* the framework is offering against adjacent positions, and *what* it is honest about not offering.

Against Markov-blanket and active-inference accounts of nested agency, AAT’s bridge-lemma form supplies a *control-theoretic loss bound* on treating a collection as one agent — something the Markov-blanket apparatus does not produce, even when the boundary criterion is in hand. The closure-defect framework is consistent with the approximate-information-state and influence-based-abstraction lines from PAC-MDP work (Subramanian 2020, Congeduti 2020, Abel 2016, Taylor 2008) but applies the same predictive-loss-to-control-error machinery at the *agent boundary formation* scope rather than at the single-agent compression scope. The framework’s distinctive content is the *use* of bounded-loss composition as an agent-boundary criterion, together with the admissibility conditions (A1)–(A4) on the macro-dynamics that force the macro-description to itself be AAT-shaped (ruling out trivial curve-fits where any composite could be called an “agent” by re-defining state).

The honest scope statement worth keeping in view is that $\varepsilon^* = 0$ does not mean *optimality* — two non-communicating Kalman filters achieve $\varepsilon^* = 0$ while being suboptimal versus a joint filter, because the joint filter could exploit cross-correlation the non-communicating pair cannot. The framework diagnoses *representability* of a composite as an AAT agent, not optimality of that composite against alternative architectures. Strong monotonicity is the hinge — sub-agents that are merely Lyapunov-stable do not compose into a coherent macro-agent, and this is one of the chapter’s quieter surprises. Three agent tiers fall out of the strong-monotonicity requirement: Tier 1 where contraction is proved for the full class (Bayesian on exponential families, gradient descent on strongly convex losses, linear correctors with positive-definite gain); Tier 2 where it holds locally; Tier 3 where it must be verified per-domain.

The critical-mass closed form unifies four signed special cases.

The chapter's deepest cross-segment finding is what `#deriv-critical-mass-composition` produces by composing the bridge lemma with `#der-team-persistence's` signed-coupling decomposition. For the symmetric-matched-Tier-1 dyad, the composite sector constant is now derived in closed form:

$$(\alpha - C)R > \rho + \gamma\mathcal{T} \quad (13.35)$$

with signed γ (negative cooperative, positive adversarial) and C the coordination-overhead cost. The single inequality recovers *four prior results as signed special cases*: single-agent persistence ($\gamma = 0$, $C = 0$, reducing to Part I Ch.4's central inequality); team-persistence with cooperative coupling ($\gamma < 0$, where the negative term enters as effective-disturbance reduction); adversarial destabilization ($\gamma > 0$, where the positive term amplifies disturbance); and coordination-dominated failure ($C > \alpha$, where the composite's effective correction rate goes negative regardless of individual quality — *Brooks's Law as a derived inequality* rather than a folkloric observation). The teleological unity U_O enters multiplicatively on γ plus scope-gate, not additively — a distinction with downstream implications for Ch.3's unity-closure-mapping result.

The implications-side payoff that the segment's Working Notes flag but is worth surfacing: this is the framework's *clean-est internal-economy demonstration*. One inequality, with signed parameters indexing the regime, recovers what previously read as four separate results. The result composes with the sector-persistence template economy (Part I Ch.4 implications, §4): the template absorbed the Lyapunov boilerplate; the critical-mass derivation specializes the template's persistence inequality with a signed-coupling effective-disturbance decomposition that handles single, cooperative, adversarial, and coordination-failure regimes simultaneously. For a practitioner, the engineering anchor is that Brooks's Law (adding people to a late project makes it later) has a *physics-level* derivation here: a development team's macro-correction rate is $\alpha - C$ where C is coordination overhead; once C exceeds α , adding people degrades persistence regardless of individual quality. The threshold is sharp, not gradual, and the inequality predicts exactly where it bites.

The closed form also produces `#disc-identifiability-floor's` Instance 3 in its fullest form. Composite contraction certification from component-marginal data is structurally impossible — the coupling-sign bit ($\gamma < 0$ vs $\gamma > 0$) is unidentifiable from how each component looks in isolation (Liberzon 2003's common-Lyapunov-nonexistence is the external no-go). The five boundary routes named in `#disc-identifiability-floor` for Instance 3 escape this: under matched-Tier conditions the closed form supplies the coupling sign by interventional access at the composite scope, via Mode 2 of the loop-Level-2 pattern (observer-on-sub-agent, from Part II Ch.2 implications). The chapter is therefore where the *third* formal instance of the identifiability-floor pattern becomes statable; the four-instance pattern is completed with Instance 1 (Part II Ch.4), Instance 2 (Part II Ch.3), Instance 3 (here), and Instance 4 (Part II Ch.5 / Appendix A — `#der-architecture-noidentifiability`, architecture-noidentifiability from on-policy summary data).

Class coercion via wrapping — what the construction unlocks.

`#der-class-coercion-via-wrapping's` Findings cover the constructive directed-separation result well: when a component (like an LLM) has belief-update and goal-conditioning entangled in a single forward pass, a scaffold around it can maintain explicit separate stores for belief and goal, with belief updates only seeing queries that don't include the goal as input. Under named admissibility conditions on the component, the wrapped system is goal-blind in its belief updates *by construction* — even though the underlying component is not. The $W_0/W_2/W_1$ regime hierarchy names the practical landscape: W_1 (strict wrapping, structural leakage bound, theoretically clean), W_2 (partial wrapping with goal-conditioned call and parsed response, behavioral leakage bound, operationally common), W_0 (no wrapping — the unconverted Class 3 case).

The implications-side payoff is what the construction unlocks beyond the segment's statement: it is the *constructive bridge* that resolves how AAT's Part II results — derived for Class 1 (Separated) agents — apply to Class 3 (Coupled) substrates. Without the wrapping construction, Class 3 components would sit outside Part II's persistence machinery

as a scope exit; with the wrapping construction, the wrapped system at the *wrapper level* is Class 1 (Separated) by construction, inherits the sector-persistence template, and operates under the standard AAT results. The cost is paid in two places: a Brooks's-Law tempo overhead (more component calls per macro-step), and a residual leakage rate bounded structurally (W_1 , by the wrapper's query design) or behaviorally (W_2 , by the component's instruction-following fidelity).

The composition-level consequences are derived in *#der-class-coercion-in-composition*: the wrapped system is a valid AAT composite agent, satisfying admissibility conditions (A1)–(A4) of *#form-composition-closure* under wrapper-design constraints (D-A2)–(D-A4). It inherits the sector-persistence template at the wrapper level; its tempo cost is the Brooks's-Law composition cost from *#der-tempo-composition*. This is the chapter's quiet load-bearing claim for *03-llm-core/*: language-model substrates are *constructively* in AAT scope, at a measurable cost, rather than scope-out. The implications for AI-safety practice are concrete: scaffolded loops around LLMs are not just engineering convenience — they are the structural mechanism that converts a Class 3 (Coupled) substrate into a Class 1 (Separated) composite, and the leakage bounds tell the practitioner what residual coupling remains. The Generative Agents observation→memory step (Park et al. 2023) is the canonical empirical W_1 instance; most production tool-using LLM frameworks (ReAct, Reflexion, Toolformer, MemGPT) sit in W_2 . The framework's regime hierarchy gives all of these constructions a *structural* reading rather than an engineering-pattern reading.

The cross-domain transfer follows naturally. Human bicameral / dual-process models (Kahneman 2011) are W_1 -style internal scaffolding when the deliberative system constrains the intuitive system through structural commitment to goal-blind belief-update queries; organizations where executive scaffolding (boards, committees, decision rights) recovers cascade ordering across coupled individuals are W_1/W_2 at the organizational scale. The framework's structural claim is the same in all three cases: *separation at the wrapper level is what gives AAT-grade behavior; the wrapping construction is structurally non-optional for Class 3 (Coupled) substrates.*

Tempo composition, Brooks's Law, and the bandwidth connection.

#der-tempo-composition carries the sub-additive tempo inequality: $\mathcal{T}_c \leq \sum_i \mathcal{T}_i$, with the gap bounded below by $C_{\text{coord}} \geq \varepsilon^* \nu_c$. The implication is that *coordination overhead acts as effective disturbance at the composite level* — closure defect $\varepsilon^* \nu_c$ is what the template's instantiation at the composite identifies as the internal contribution to effective disturbance, alongside the external environmental rate ρ_{ext} . The composite persists when $\alpha_c R_c > \rho_{\text{ext}} + \varepsilon^* \nu_c$, and Brooks's Law follows: adding agents increases $\sum_i \mathcal{T}_i$ but may increase $\varepsilon^* \nu_c$ faster, pushing the composite's effective disturbance past its persistence threshold even as nominal capacity grows.

The composition with Part I Ch.4's information-rate floor is worth surfacing as a cross-segment finding. Each agent in the composite pays its own $\dot{R} \geq n\alpha/2$ bandwidth cost; the composite as a whole pays the sum *minus what coordinated sub-agents save on shared observation*. But there is a complementary cost: the closure defect $\varepsilon^* \nu_c$ has its own information-rate signature, because maintaining the bridge-lemma's strong-monotonicity condition at the macro-update requires the composite to acquire information about how its sub-agents are interacting. The framework's prescription is that *composite persistence pays for both the sub-agents' individual bandwidth and the coordination-overhead bandwidth*, and the trade-off between adding agents (more $\sum_i \mathcal{T}_i$, more individual bandwidth) and increasing coordination overhead (more $\varepsilon^* \nu_c$, more coordination bandwidth) has a precise inflection point at $C_{\text{coord}} = \alpha_c$. Past the inflection point, additional sub-agents *cost more bandwidth than they buy*, and the composite's persistence-information-rate floor moves up faster than its persistence threshold. This is the implication of the chapter's machinery for distributed-systems design: the sweet spot for sub-agent count is bounded by the coordination-overhead growth rate, not by individual sub-agent capacity.

The chapter is also where Mode 2 of the loop-Level-2 deployment pattern (observer-on-sub-agent, from Part II Ch.2 implications) becomes structurally load-bearing. The Liberzon-anchored identifiability-floor Instance 3 says the coupling sign cannot be recovered from component marginals; the unique broadly-available escape is interventional access at the composite scope. Operationally, this means an external observer (or the composite agent itself acting on its sub-agents) can recover the sign by performing *do*-interventions on one sub-agent and observing the response of another.

The closed-form critical-mass derivation gives the mathematical form of the recovered coupling; the loop-Level-2 access at the composite level gives the data substrate that recovery requires.

Closing notes and the bridge into Unity, Communication, and Shared Intent.

What this chapter closes on is the composite-agent foundation Part III needs: a representability criterion via closure defect, a trajectory-level loss bound via the bridge lemma, a closed-form persistence inequality that recovers single-agent, team, adversarial, and Brooks's-Law regimes as signed special cases, a constructive route from Class 3 (Coupled) components to Class 1 (Separated) composites with measured leakage bounds, and a sub-additive tempo inequality whose coordination-overhead term enters the persistence template as effective disturbance at the composite level. The chapter's two segment-level ## Findings sections carry the load-bearing technical content; this implications segment surfaces the cross-segment composition and the engineering anchors that the segments individually cannot reach.

The bridge into Part III Ch.3 (Unity, Communication, and Shared Intent) is the closure-defect ε^* itself. The dimensions of unity — content axes ($U_M, U_O, U_\Sigma, U_{\text{obs}}$) and structural axis (U_f) — parametrize *what sub-agents can share*; Ch.3's #result-unity-closure-mapping shows that unity parametrizes a rate-distortion curve for closure defect. The chapter here gives the closure defect a precise form; Ch.3 gives unity its quantitative content; together they say *which dimensions of unity matter how much for which kinds of composite-agent persistence*. The cooperative-vs-adversarial machinery in Ch.4 then operates on what Ch.2 and Ch.3 jointly build: cooperative coupling reduces γ (lowers effective disturbance); adversarial coupling raises γ (amplifies it); the signed-coupling structure is the same inequality viewed from opposite directions. Ch.5's strategic composition lifts the symmetric partially-opposing case from scalar coupling to equilibrium machinery (potential / monotone games), and the identifiability-floor pattern's Instance 3 instance becomes operational there in its strongest form.

Working Notes.

- This segment is a chapter-end discussion grouping findings from #form-composition-closure, #der-tempo-composition, #hyp-directed-separation-under-composition, #der-class-coercion-via-wrapping, #der-class-coercion-in-composition, #deriv-critical-mass-composition (Appendix A), #result-sector-persistence-template (Appendix A), and #der-directed-separation (Part II Ch.1, the architectural classification the wrapping construction lifts to composite status). Cross-pattern reading leans on #disc-identifiability-floor's meta-segment for the Instance 3 surfacing and #disc-modularity-state-dynamics for the truthification-mechanism reading of the wrapping construction.
- **Cross-reference vs canonical-home policy.** Two source segments in this chapter carry well-written ## Findings sections — #form-composition-closure (closure-defect and bridge lemma) and #der-class-coercion-via-wrapping (constructive directed separation). Both are cross-referenced rather than restated. The findings whose canonical home is this implications segment are the *cross-segment compositions*: the critical-mass closed form as recovery of four signed special cases (composing #deriv-critical-mass-composition with #der-team-persistence and Part I Ch.4's persistence condition), the Brooks's-Law derivation as inequality (composing #der-tempo-composition with the template), the composition-level identifiability-floor Instance 3 (composing #deriv-critical-mass-composition with #disc-identifiability-floor's pattern and Part II Ch.2's loop-Level-2 access), and the wrapping-construction-as-truthification (composing #der-class-coercion-via-wrapping with #disc-modularity-state-dynamics's three-operation pattern).
- **Class-coercion as truthification mechanism — connection to #disc-modularity-state-dynamics.** The M4 meta-segment positions modularity as a *contested property* under three operations: truthification (self-driven-increasing), strategic self-coupling (self-driven-decreasing), and adversarial coupling pressure (externally-driven-decreasing). Class coercion via wrapping is one of two operational mechanisms for truthification — alongside defensive scaffolding (peer review, prediction registers, etc., from #disc-adversarial-coupling-pressure). The chapter is therefore where modularity-as-state lands its first construction: wrapping is a structural commitment that *moves* a composite agent's modularity state from W_0 (no wrapping) toward W_1 (strict wrapping) under explicit cost, and the cost is the framework's quantitative content for the M4 truthification operation.

- **Identifiability-floor pattern coverage now four-of-four.** With this chapter’s content, the pattern’s four formal instances are all in view: Instance 1 (Part II Ch.4 — on-policy L0-insufficiency), Instance 2 (Part II Ch.3 — L1’ mixture under unobservable common cause), Instance 3 (here — composite contraction certification from component marginals), Instance 4 (Part II Ch.5 / Appendix A — architecture-noidentifiability from on-policy summary data, #der-architecture-noidentifiability). Instance 5 (Mehra non-identifiability in adaptive Kalman) remains candidate. The implications-side payoff is the cross-instance dual structure: Instance 1 and Instance 2 cover agent-internal failures of identification (single agent, action / mixture layers); Instance 3 covers composition-level failure (coupling sign across sub-agents); Instance 4 covers architecture-level failure (similarity orbit on the agent’s internal realization). Three of the four (Instances 1, 2, 4) sit in the rank-collapse subclass with Sylvester’s law of inertia governing escape-irreducibility (directly for Instances 1 and 2; at one remove for Instance 4, whose generating action is state-space similarity rather than metric congruence); Instance 3 is the distinct projection/closure member. The unique broadly-available escapes share Pearl-do structure (loop-interventional access) — Instances 1, 3, and 4 all elevate #der-loop-interventional-access as *three semantically-distinct deployment modes* (agent-self-intervention / observer-on-sub-agent / observer-on-agent-input), each breaking a distinct Level-1 degeneracy.
- **No NeurIPS submission directly formalizes this chapter’s machinery.** The closure-defect framework, the bridge lemma, and the critical-mass closed form are AAT-internal; their cross-references in Paper 2 (RL convergence) and Paper 3 (LLM hallucinate bound) are scope and structural-argument links rather than theorem-grade adoption. A composition-themed submission would naturally land here in a future cycle.

Incidental audit gold (lift 2026-05-30). Cross-audit “wandering thoughts” / §14-ideation harvested from the de-novo auditors’ working dirs. **Coverage:** only one dir reached a digested reflection on this chapter-end implications segment (526815) — the deeply-mathematical critique pass. **No pedagogical / analogy / figure gold surfaced** — 526815’s reflection is entirely *certified-track correctness findings* (the F139–F145 cluster, mostly “this implications map presents appendix-derived / conditional payoffs as settled”), routed to the findings track, not lifted as Working-Notes gold. The non-finding signal worth preserving for the eventual synthesis:

Belongs elsewhere / convergence signal.

- **The representability-vs-optimality distinction was named the chapter’s strongest contribution to preserve in the final synthesis** — “the representability-versus-optimality distinction is one of the chapter’s strongest contributions. Preserve it in final synthesis” (Codex/Claude, AUDIT-WORKING-526815, “Watch”). Converges with the same distinction lifted into 13.1 (the two-Kalman result) — a clear signal this is the chapter’s headline idea to foreground in the implications map’s framing.

Off-ramp note: 526815’s main thrust here is that this discussion-grade implications map “presents the critical-mass closed form, four signed special cases, identifiability-floor Instance 3, and bandwidth-inflection claims as settled” while their proof homes are appendix-derived or conditional, and that it inherits the upstream tempo-double-counting and wrapper-scope-gate issues (F139–F143). These are voice-discipline / scope-honesty matters for the implications segment, **not** gold — flagged in the lift report for the findings/adjudication track.

Chapter 14

Unity, Communication, and Shared Intent

[Gap] Discussion

The quality of a composite agent's composition — *conditional on 12.2 being satisfied via at least one of its four routes (three alignment routes + the strategic-equilibrium route (C-iv))* — is parametrized along **two architecturally distinct axes**:

- **Content axis (four unity dimensions).** What the sub-agents *share*: epistemic (U_M , shared model), teleological (U_O , shared objective), strategic (U_Σ , coordinated action), and perceptual (U_{obs} , shared observations).
- **Structural axis (update-rule homogeneity, U_f).** Whether sub-agents implement the *same* correction rule: how similar their f_M updates are across the population.

Together, the two axes parametrize the rate-distortion curves for the component closure defects (13.1, 14.1). Higher unity along either axis permits more aggressive compression at lower closure defect; neither axis alone is sufficient. In pure Part I composition (passive estimators, no G_t), agents with identical content can still produce non-zero ε_x if their update rules differ — the content axis cannot detect this, which is why the structural axis is required. Unity (in either sense) does not directly predict closure-defect magnitude; it controls the compressibility of the corresponding state, observation, or action component under projection.

Scope. The decomposition applies to composites that satisfy 12.2. U_O plays a role in the (C-i) route of the scope condition via value-correlation, but the scope condition is a disjunction of four routes — three alignment routes (shared objective, hierarchical derivation, mutual benefit) plus the strategic-equilibrium route (C-iv) — not a scalar threshold on U_O . Below scope-satisfaction (no route applies), the sub-agents are a multi-agent system per 12.1 and composition-level quantities are not well-defined.

For a composite agent A_c composed of sub-agents $\{A_1, \dots, A_n\}$, the unity profile consists of *four content dimensions* (this section, below) and *one structural dimension* (U_f , defined after the content dimensions).

Definition (definition-unity-dimensions).

Epistemic unity U_M — how much of the reality model is shared:

$$U_M = \frac{I(M_t^{(1)}; \dots; M_t^{(n)})}{(n-1) H(M_t^{(1)}, \dots, M_t^{(n)})} \quad (14.1)$$

The fraction of total model information that is shared (total-correlation ratio, normalized to $[0, 1]$). The numerator is the total correlation $I = \sum_i H(M_t^{(i)}) - H(M_t^{(1)}, \dots, M_t^{(n)})$; the $(n - 1)$ factor in the denominator is its normalizer, since each marginal entropy is at most the joint entropy and hence $I \leq (n - 1) H(M_t^{(1)}, \dots, M_t^{(n)})$, with equality iff each model determines all the others. So $U_M = 1$ for identical models and $U_M = 0$ for independent models, for every n . (Without the $(n - 1)$ factor the ratio reaches $n - 1$ at identity — the bare multi-information/joint-entropy form is normalized correctly only at $n = 2$.)

Teleological unity U_O — how aligned are the objectives:

$$U_O^{(i,j)} = \text{corr}\left(V_{O_t^{(i)}}(\tau), V_{O_t^{(j)}}(\tau)\right) \quad (14.2)$$

over trajectories the composite encounters. $+1$ for identical objectives; -1 for perfectly opposed; 0 for orthogonal. The composite teleological unity is an aggregation over all pairs. The scalar ranges from fully cooperative to fully adversarial per objective dimension.

[Scope note] U_O plays a role in 12.2 — primarily along route (C-i), where value-correlation is the operationalization of teleological alignment. The scope condition itself is disjunctive: it is satisfied when *any* of the alignment routes (C-i), (C-ii), (C-iii) applies, or when the strategic-equilibrium route (C-iv) applies; not by U_O alone crossing a common scalar threshold. The quality-metric role of U_O captured in this segment presumes scope-satisfaction via *some* route; on the (C-iv) side it tracks at best the alignment projection of a strategic composite whose macro-state is defined by equilibrium structure rather than by a shared target. When the sub-agents fail all four routes, they form a multi-agent system (12.1) but not a composite, and composition-level quantities (closure defect, composite tempo, team persistence) are not well-defined.

Strategic unity U_Σ — how coordinated is the joint policy:

Discussion.

$$U_\Sigma = 1 - \frac{D_{\text{KL}}(\pi_{\text{actual}}^c \| \pi_{\text{optimal}}^c)}{D_{\text{KL}}(\pi_{\text{independent}}^c \| \pi_{\text{optimal}}^c)} \quad (14.3)$$

where π_{optimal}^c is the jointly optimal policy. $U_\Sigma = 1$ when actual matches optimal; $U_\Sigma = 0$ when actual matches independent (no coordination). Requires knowing the jointly optimal policy, which is itself a strong assumption.

Perceptual unity U_{obs} — how much of the observation stream is shared:

The fraction of total observation information that reaches all sub-agents. Full perceptual unity means all agents observe the same signals; zero means private observations only. Enables epistemic convergence without explicit model-sharing.

Update-rule homogeneity U_f — how similar the sub-agent update rules are:

Definition.

$$U_f = 1 - d\left(f_M^{(1)}, \dots, f_M^{(n)}\right) \quad (14.4)$$

where d is a distance over the space of update operators $f_M : (M, o, a) \mapsto M'$, normalized so that $U_f = 1$ when all sub-agents implement the same correction rule and $U_f = 0$ at maximal heterogeneity. The choice of d is case-specific — for parametric Kalman-like updates, $d \propto |\Delta K|/K_{\text{max}}$ on the gain parameter; for Bayesian updates with shared structural form but different priors, d tracks divergence between the induced kernels; for arbitrary f_M in function space, candidates include operator-norm distance, Fisher-information-weighted distance, or IB-style comparison.

Where the four content unities measure shared *information* across sub-agents (state, objective, policy, observation), U_f measures shared *structure* — whether the agents instantiate the same update law. The two axes are independent: agents can share a model ($U_M = 1$) while updating it differently ($U_f < 1$), and conversely. In purposeful settings (G_t present), U_Σ partially absorbs structural variation in the policy half of the cycle, but the model-update half remains uncovered without U_f . In pure Part I composition (passive estimators, no G_t), U_f is the only handle on structural homogeneity. The closed-form linear-Gaussian instance — heterogeneous Kalman gains $\Delta K = K_1^* - K_2^*$ producing $\varepsilon_x \propto |\Delta K|$ — is derived in 14.1 §Two-axis structure.

Discussion.

The achievable component closure defect $\varepsilon_d^{\min}(k_d)$ under a projection of macro-dimension k_d is a function of *both* axes — the relevant content unity U_d and the structural unity U_f — together with the projection-dimension parameter:

$$\varepsilon_d^{\min}(k_d) = f_d(k_d; U_d, U_f) \quad (14.5)$$

monotone decreasing in each unity argument and monotone increasing in compression aggressiveness (smaller k_d). The form is derived in 14.1; in linear-Gaussian scalar cases it admits closed-form expressions for $d \in \{x, o, a\}$. U_O and U_Σ enter ε_a jointly rather than separately — collapsing alignment to a single scalar (U_O alone) would mis-route the repair, since execution-path divergence (U_Σ) cannot be fixed by re-aligning targets. This is the **anti-collapse** discipline (5.7) at the composite layer: two distinct quantities a naive reading merges, routing to different corrections.

Discussion-grade. Max attainable: empirical. The four content dimensions are qualitatively motivated by correspondence with the four components of agent state (M_t , O_t , Σ_t , and the observation channel); the structural dimension U_f is forced by the linear-Gaussian two-Kalman case (14.1 §Two-axis structure), where heterogeneous gains produce non-zero ε_x that no content dimension can register. The two-axis architecture (content $\times$ structure) is therefore a definitional commitment of this segment, not a heuristic.

The specific metrics are sketches. The information-theoretic formulations (U_M , U_Σ) are well-defined in principle but require specifying distributions and distance measures for practical computation. U_f is even less prescriptive — the choice of operator distance d is case-specific, and a general theory of structural-variation measures across arbitrary f_M classes is open.

The claim that the dimensions are *substantially independent* is a hypothesis, not derived. Epistemic unity may enable strategic unity (shared models allow coordination without explicit planning); content unity along U_M does not constrain U_f (agents can share a posterior while updating it differently). Independence holds approximately, with documented joint dependencies: (U_O , U_Σ) jointly control ε_a and cannot be separated; U_f enters all three closure components (ε_x , ε_o , ε_a) as a structural multiplier on what content unity alone would predict.

Clausewitz’s three gaps. These dimensions map to the gaps identified by Clausewitz (systematized by Bungay in *The Art of Action*):

Table 14.1

<i>Clausewitz Gap</i>	<i>Unity Dimension</i>	<i>Formal Quantity</i>
Knowledge gap	Epistemic unity (U_M)	$1 - U_M$: fraction of model not shared
Alignment gap	Teleological unity (U_O)	$1 - U_O$: objective misalignment
Effects gap	Strategic + Perceptual unity	$1 - U_\Sigma$ + observation routing costs

The mapping is not perfect — Clausewitz’s “effects gap” blends action coordination with observation feedback — but it provides 200+ years of organizational evidence for the qualitative decomposition.

Connection to closure defect. The unity dimensions parametrize a rate-distortion relation with the component closure errors in 13.1, not a direct correspondence. The formal statement is in 14.1: the achievable closure-defect component $\varepsilon_d(k_d)$ under a projection of macro-dimension k_d decreases monotonically in both the relevant content unity U_d and the structural unity U_f , with closed-form expressions in the linear-Gaussian case. Qualitative direction along the content axis: U_M governs the compressibility of state information (ε_x), (U_O, U_Σ) jointly govern action compressibility (ε_a), U_{obs} governs observation compressibility (ε_o). The naive reading “high U_d predicts low ε_d ” fails: for non-compressing projections (e.g., the means-only projection in the two-Kalman case) $\varepsilon_x = 0$ regardless of U_M , while for heterogeneous-gain projections $\varepsilon_x > 0$ even at perfect content correlation. Both observations point at the same correction — closure defect lives on a rate-distortion surface with two unity arguments, not a single one.

What each dimension’s absence costs.

- Low U_M : prediction conflicts $\rightarrow$ uncoordinated actions based on contradictory beliefs. Internal mismatch component from model disagreement.
- Low U_O : strategic friction $\rightarrow$ sub-agents pursue conflicting sub-goals. Effort wasted or counterproductive.
- Low U_Σ : redundancy and gaps $\rightarrow$ two agents fix the same bug while a critical one goes unnoticed.
- Low U_{obs} : information silos $\rightarrow$ critical signals observed by one agent but not actionable by the composite.
- Low U_f : structural drift $\rightarrow$ even agents sharing identical models, objectives, observations, and policy targets produce divergent macro-trajectories under aggressive projection, because their corrections respond to the same evidence with different gains. The composite cannot be summarized at low macro-dimension without residual error scaling with the gain mismatch.
- The independence of unity dimensions needs careful examination. High epistemic unity likely enables (but does not guarantee) high strategic unity — if agents share models, they can coordinate implicitly. The dimensions may be better described as “substantially independent inputs to a joint prediction of ε^* ” rather than “independent properties.” Independence between content axis and structural axis ($U_d \perp U_f$) is cleaner — content sharing and update-rule similarity are categorically distinct properties — but a formal proof is open.
- The specific metric formulations need testing on concrete cases (software team, military unit) to determine if they discriminate meaningfully between well-composed and poorly-composed groups.
- The teleological unity scalar per objective dimension (+1 to -1) captures mixed cooperative-competitive situations: a company can be cooperative on product quality and competitive on internal resource allocation simultaneously.
- **U_f operator-distance choice is open.** The definition leaves d case-specific. A general theory of structural-variation measures across arbitrary f_M classes (operator-norm distance, Fisher-information-weighted distance, IB-style comparison) is unsettled. The linear-Gaussian Kalman case ($d \propto |\Delta K|/K_{\text{max}}$) is the only worked closed form; non-Gaussian and non-linear cases are open follow-up work tracked in 14.1 Working Notes.
- **Joint (U_O, U_Σ) $\rightarrow \varepsilon_a$ dependence.** State error tracks U_M ; action error tracks *both* U_O (target alignment) *and* U_Σ (policy alignment); observation error tracks U_{obs} . The two dimensions jointly controlling action error are physically distinct: U_O is about evaluation/preference agreement; U_Σ is about execution-path agreement. Agents with identical objectives but different execution plans have high U_O , low U_Σ ; agents coordinating on arbitrary shared routines have high U_Σ , low U_O . See 14.1 for the quantitative relationship.

- **U_O as scope vs. quality.** The scope role is in 12.2 as a disjunction of four routes (three alignment + one strategic-equilibrium) — U_O is the value-correlation operationalization of route (C-i), not a universal scope variable. The quality-metric role remains here, presumed conditional on scope-satisfaction via at least one route. Open: whether the three alignment routes (C-i)–(C-iii) reduce to a single scalar is not established, and the strategic-equilibrium route (C-iv) is structurally distinct from any alignment-side aggregation, so “scope-satisfaction” and “ U_O value” should be treated as distinct concerns in downstream uses.

Cross-audit “wandering thoughts” / §14-ideation harvested from the de-novo auditors’ working dirs, deduplicated across substrates and attributed by substrate + audit. Orthogonal pedagogical / framing / figure / naming material, kept separate from the certified theory-fix findings (the metric-repair findings F146–F150 from AUDIT-WORKING-526815 are routed for adjudication, not here — see the off-ramp note at the end). **Coverage:** three dirs carry a dedicated reflection (526815, 829314, 849201) plus two batched dirs (451729 batch-14, 613842 composition-cross-component). The other contributing dirs stopped earlier in the OUTLINE walk. Substrate attribution inferred from voice where not explicit; uncertain cases hedged. **Finding-vs-framing conflation preserved:** several “candidate Discussion” entries below were written as confident results in first-encounter cognition; that texture is signal, kept rather than sanitized.

- Plain-statement of the structural-axis insight, Feynman-criterion-adjacent: “even if two agents know the exact same things and want the exact same things, if one learns fast and the other learns slow, they cannot be mathematically compressed into a single macro-agent without accumulating trajectory error” (Gemini, AUDIT-WORKING-829314). Compact hook for a Brief or Discussion opener on why content unity is not enough.
- “Decomposes ‘teamwork’ into its orthogonal mathematical components” — one-line gloss of the segment’s contribution (Claude, AUDIT-WORKING-849201).
- **Clausewitz-three-gaps mapping is the standout pedagogy — promote it more prominently.** Two substrates independently flagged the mapping of the unity dimensions to Clausewitz’s three gaps (Knowledge / Alignment / Effects) as “an absolutely brilliant piece of pedagogical framing ... it grounds the abstract information theory in 200 years of practical military strategy”; the recommendation is to highlight that table even more prominently (Gemini, AUDIT-WORKING-829314; Gemini, AUDIT-WORKING-849201 — “a phenomenal domain connection”).
- **Content alignment is dynamically unstable without structural alignment (the “metabolism mismatch” / “culture fit” framing).** The sharpest worked elaboration of U_f : a two-engineer team perfectly aligned on content ($U_M = U_O = U_\Sigma = 1$) but with mismatched update rules — Engineer A highly plastic (rewrites architecture on every bug), Engineer B highly stable (patches the line) — will violently diverge the instant a shared bug is observed, “not because they disagreed on facts or goals, but simply because their internal learning metabolisms were mismatched.” This is offered as the formal content of the corporate “culture fit” problem: hiring is for f_M -compatibility, not only skills; dropping a “move fast and break things” developer into a slow-moving aerospace team spikes the closure defect ε^* and destroys macro-agent tempo (Gemini, AUDIT-WORKING-829314). (*Finding-vs-framing texture: written as a proof — “proves mathematically that ...” — but the general- U_f monotonicity is only worked for the linear-Gaussian two-Kalman case per 14.1; promote as motivated framing, not as established theorem.*)
- **Epistemology does not guarantee strategy — the two-autonomous-cars example.** “Two autonomous cars might have perfect, identical models of an intersection ($U_M = 1$), but if they don’t coordinate their policies, they will still crash ($U_\Sigma = 0$).” A vivid one-line separation of the epistemic and strategic content axes (Gemini, AUDIT-WORKING-829314).
- **Closure defect as the formal analog of “organizational complexity.”** ε^* reads as “the gap between what a collection of agents could achieve in principle and what a coherent macro-description can represent” — the formal version of what organizational scientists call organizational complexity. Crucially: coherence (closure

defect) and efficiency (performance gap) are *orthogonal* axes — “you can have a perfectly coherent organizational description of an inefficient organization” (Claude, AUDIT-WORKING-451729, batch-14). A candidate framing-level Discussion note on what the two-axis taxonomy is *for*.

- **The two-Kalman $\epsilon^* = 0$ diagnostic, in words.** Two non-communicating Kalman filters tracking correlated targets are “perfectly representable as one agent” even though they never share information — not because they cooperate, but because each individually generates a sufficient statistic for the macro-state’s next observation. Representability $\neq$ optimality (Claude, AUDIT-WORKING-451729, batch-14). Belongs more to 14.1 but motivates why the content axis alone is the wrong diagnostic here.
- **Mixed-motive U_O .** How does the theory handle U_O in mixed-motive games where agents cooperate on survival but compete on resource allocation? The per-objective-dimension scalar (+1 to −1) is noted to cover this in principle, but the aggregation math “must be tricky” — worth a worked example (Claude, AUDIT-WORKING-849201; echoed Gemini, AUDIT-WORKING-829314 on group-level/nontransitive alignment).
- **C-iv ontology drift — closure/unity machinery still assumes the older shared-objective idiom.** A cross-component pass flagged that the scope layer (12.2) has accepted route (C-iv) (strategic composites with no shared O_c , macro-state defined relative to equilibrium structure $\mathcal{E}$), but this segment “still describes the scope as a three-route disjunction and treats composition-level quantities as unavailable whenever the three alignment routes fail” — a real integration gap between the broadened scope and the unity machinery (Claude, AUDIT-WORKING-613842). The existing “ U_O as scope vs. quality” Working Note above already partly tracks this; the cross-component observation sharpens it into a propagate-C-iv-into-the-unity-segments task.
- Auditors converged on the same conceptual checks as natural reader questions: that $U_M = 1$ does not force $U_\Sigma = 1$ (the cars), and that the four content unities measure shared *content* but not shared *structure* — the “missing axis” that U_f supplies. Two fresh readers independently praised the segment for flagging this gap rather than hiding it (“the epistemic honesty to flag this gap ... is exemplary” — Claude, AUDIT-WORKING-849201; Gemini, AUDIT-WORKING-829314).
- **Two-axis closure-defect surface.** A rate-distortion surface controlled by two roughly-independent inputs — content unity and structural homogeneity — with the metric-normalization warning marked on the content side (since bad normalization would mislead the surface before any later derivation begins) (Codex/Claude, AUDIT-WORKING-526815).
- **Rate-distortion-as-compressibility army/management framing** belongs to 14.1 (lifted there), but originates partly from the structural-axis intuition seeded here.
- AUDIT-WORKING-526815 raised five metric-level findings on this segment that are strengthen-first / repair candidates, not framing gold, and are flagged here only so they are not lost: **F146** — $U_M = I(M^{(1)}; \dots; M^{(n)})/H(M^{(1)}, \dots, M^{(n)})$ is *not* normalized to $[0, 1]$ as the discussion-grade label implies (for n identical variables of entropy H , total correlation is $(n - 1)H$ and joint entropy is H , so the ratio is $n - 1$, not 1 — needs a capped/rescaled redundancy measure); **F147** — teleological unity as pairwise value-function correlation is distribution-/policy-dependent and undefined for zero-variance value functions (needs support/variance/aggregation conventions); **F148** — the strategic-unity KL ratio can have a zero denominator, infinite KL, or leave $[0, 1]$ without clipping; **F149** (soft) — U_{obs} as “fraction of observation information reaching all sub-agents” misses complementary-private-plus-routed and synergistic-private observation; **F150** (soft) — $U_f = 1 - d(\cdot)$ presumes a normalized $d \in [0, 1]$ that the candidate operator/Fisher/IB distances do not automatically satisfy. *These are metric-formalization re-pairs (mostly scope/normalization tightenings), not no-gos; routed for adjudication on the strengthen-first track.*

14.1 Result: Unity-to-Closure Rate-Distortion Mapping

Unity dimensions parametrize rate-distortion curves for closure-defect components, not point-valued predictors. The achievable closure-defect component ε_d under projection of macro-dimension k_d is monotone decreasing in both the relevant content unity U_d and the structural unity U_f (update-rule homogeneity), with higher unity along either axis lowering achievable defect at a given compression. Closed forms hold in the linear-Gaussian case; structural monotonicity survives more broadly. The two-axis structure (content $\times$ structure) is forced by the heterogeneous-Kalman case below and is reflected definitionally in `#def-unity-dimensions`.

Several closed-form linear-Gaussian instances drive the structural conclusions. *Observation closure defect* scales as $1 - U_{\text{obs}}$ in the two-agent scalar case under 1D projection: higher perceptual unity means observations are more redundant, a 1D summary suffices, and ε_o shrinks. *Action closure defect* scales as $\kappa^2(1 - U_O)$ for independent policies under 1D state projection, with policy coordination (strategic unity) providing a further multiplicative reduction — U_O and U_Σ enter ε_a *jointly* rather than separately. A striking structural fact for the *state* closure defect: under linear dynamics with consistent linear projections, $\varepsilon_x \equiv 0$ regardless of U_M or compression dimension. State closure becomes non-trivial only when the projection's range is non-invariant under the micro-dynamics, the projection is inconsistent, the micro-dynamics are nonlinear — or *sub-agent update rules are heterogeneous*. This is the structural fact that forces the two-axis decomposition: the content axis cannot detect structural heterogeneity in updates; the structural axis is required.

The heterogeneous-Kalman closed form makes this concrete: two Kalman filters with different gains produce $\varepsilon_x^2 \propto (\Delta K/2)^2$ — a defect proportional to the gain mismatch even when priors are identical, even at perfect content correlation. The same structural-unity mechanism lifts to the *strategy layer* via the credence-composition case: N agents reasoning over a shared plan skeleton, working in the log-odds coordinate forced by `#deriv-edge-update-natural-parameter`, produce ε_Σ^{*2} driven by the *variance of edge-update gains across agents* — homogeneous gains yield zero defect, gain dispersion drives the defect, and the result is *dimension-free* in N . The structural-unity axis is forced — not optional — for any framework that wants to predict closure defect from agent-level properties.

The conceptual consequence of the rate-distortion framing is that **unity is a compressibility parameter**, not a magnitude that maps to closure-defect magnitude. Higher unity allows more aggressive compression at the same closure defect, or equivalently, lower closure defect at the same compression. The rate-distortion *surface* is what unity controls; closure-defect values sit on that surface depending on how aggressively the composite is compressed. This is the same shape as Shannon rate-distortion theory and as the Information Bottleneck (`#form-information-bottleneck`) — composition trades representational compactness against fidelity-to-the-micro-system, and unity dimensions name where the trade is cheap (high unity, small defect achievable under aggressive compression) versus expensive (low unity, defect grows fast under any non-trivial compression).

Formal Expression.

Rate-distortion framing (general).

Fix a composite agent satisfying the admissibility conditions (A1)-(A4) in 13.1. For each content unity dimension U_d (with $d \in \{M, \Sigma, \text{obs}\}$, and U_O contributing jointly with U_Σ — see below) and the structural unity U_f (update-rule homogeneity, defined in 14), the achievable component closure defect under a projection whose corresponding macro-dimension is k_d satisfies:

Formulation (unity-rate-distortion).

$$\varepsilon_d^{\min}(k_d) = f_d(k_d; U_d, U_f) \quad (14.6)$$

where f_d is monotone decreasing in both unity arguments, monotone increasing in aggressiveness of compression (smaller k_d). The mapping from unity to closure-defect *magnitude* is via the shape of this rate-distortion surface; unity

does not directly predict closure-defect value. In the linear-Gaussian Kalman case the structural argument reduces to $1 - U_f \propto |\Delta K|/K_{\max}$ on the gain mismatch.

Linear-Gaussian closed forms (two-agent scalar case).

For two agents with scalar observations correlated at $\rho_{o,\text{eff}}$ (combining ρ_{env} and ρ_{obs}), under 1D principal-component projection of observations, the minimum achievable observation closure defect is:

Derived (obs-closure-linear-Gaussian, from unity-dimensions, composition-closure).

$$\varepsilon_o^2(k_o = 1) = \sigma_o^2 \cdot \frac{1 - \rho_{o,\text{eff}}}{2} \propto 1 - U_{\text{obs}} \quad (14.7)$$

Higher perceptual unity $\rightarrow$ observations are more redundant $\rightarrow$ 1D summary suffices $\rightarrow \varepsilon_o$ small. Exact in the linear-Gaussian scalar case.

For two agents with scalar quadratic objectives, independent LQR policies ($\rho_\Sigma = 0$), scalar targets r_1, r_2 with correlation ρ_O , under 1D state projection, the minimum achievable action closure defect is:

Derived (action-closure-independent-policies, from unity-dimensions, composition-closure).

$$\varepsilon_a^2 \propto \kappa^2 \cdot (1 - \rho_O) \propto \kappa^2 \cdot (1 - U_O) \quad (14.8)$$

where κ is the scalar LQR gain. Adding policy coordination ($\rho_\Sigma > 0$) further reduces ε_a through a multiplicative factor. The joint (U_O, U_Σ) dependence takes the form:

$$\varepsilon_a^2 \propto (1 - U_O) \cdot f_1(U_\Sigma) + g(U_\Sigma) \quad (14.9)$$

with f_1 decreasing in U_Σ and g capturing residual strategic-misalignment error even when targets coincide.

State closure in linear-Gaussian.

For linear-Gaussian micro-dynamics with consistent linear projections Λ_x and Λ_o (the macro observation projection is the same linear combination as the macro state projection), the state closure defect vanishes:

Derived (state-closure-linear-Gaussian-trivial, from composition-closure).

$$\varepsilon_x = 0 \quad (14.10)$$

regardless of U_M or compression dimension. Linear projections of linear dynamics are exact *when the range of Λ_x is invariant under the micro-dynamics matrix*. ε_x becomes non-trivial when:

- the projection's range is non-invariant under the dynamics matrix — even with linear dynamics, consistent projections, and homogeneous updates, cross-coordinate coupling or anisotropic noise scales that mix macro-subspace components with their orthogonal complement give $\varepsilon_x > 0$ (the Mori-Zwanzig zero-lag bound $\varepsilon^* \geq \|Q_\Lambda U P_\Lambda\|_{\text{op}}$ in 13.1 is the general expression of this obstruction),
- the projection is inconsistent (macro state and macro observation projections disagree),
- the micro-dynamics are nonlinear, or
- sub-agent update rules are heterogeneous (see Two-axis structure below).

Two-axis structure (update heterogeneity).

In the non-degenerate linear-Gaussian case with heterogeneous sub-agent update rules — e.g., two Kalman filters with different gains $K_1^* \neq K_2^*$ tracking correlated processes, projected to the 1D sum $\hat{w}_+ = (\hat{w}_1 + \hat{w}_2)/\sqrt{2}$ — the state closure defect has the closed form:

Derived (two-axis-structure, from composition-closure, linear-Gaussian case).

$$\varepsilon_x^2 = (\Delta K/2)^2 [S_- - C_{+-}^2/S_+] \quad (14.11)$$

where $\Delta K = K_1^* - K_2^*$, $S_{\pm}$ are the innovation variances in the $\pm$ directions, and C_{+-} is their cross-covariance.

This exhibits two independent drivers of ε_x , one along each unity axis of 14:

1. **Content unity** (U_M , via process correlation ρ): higher correlation $\rightarrow$ lower ε_x .
2. **Structural unity** (U_f , via gain mismatch ΔK): when $\Delta K = 0$ (i.e., $U_f = 1$), $\varepsilon_x = 0$ at every ρ ; when $\Delta K \neq 0$, $\varepsilon_x > 0$ even at perfect content correlation.

The four content unities measure shared information (goals, policies, observations, model state); U_f measures whether sub-agents implement the same correction rule. The two axes contribute to the closure-defect rate-distortion surface independently — content unity controls compressibility of what the agents agree on; structural unity controls whether projection induces memory by mixing the discarded subspace into the retained one.

Strategy-layer instance (credence composition). The same structural-unity axis has an exact closed form one level up, at the strategy layer. For N agents reasoning over a shared plan skeleton (V, E) and disagreeing only on edge-credences, work in the log-odds coordinate $\lambda_{ij} = \log \frac{p_{ij}}{1-p_{ij}}$ — the unique additive-evidence coordinate (8.3; forced by U) in which edge updates are additive and the natural macro-projection is the log-odds centroid. The per-step closure defect is then driven by update-rule heterogeneity exactly as the Kalman case is, with the per-agent edge gains $\eta_{\Sigma,i}$ playing the role of the Kalman gains:

$$\varepsilon_{\Sigma}^{*2} = |E| \cdot \overline{\text{Var}_i[\eta_{\Sigma,i}]} \cdot \text{Var}[r], \quad (14.12)$$

where r is the per-edge evidence residual. This is the structural-unity (U_f) axis lifted from state to strategy: homogeneous gains ($\overline{\text{Var}_i[\eta_{\Sigma,i}]} = 0$, $U_f = 1$) give $\varepsilon_{\Sigma} = 0$; gain dispersion across agents drives the defect, scaled by plan size $|E|$ and evidence-residual variance. It is *dimension-free in N* — a population variance is estimated more precisely, not enlarged, by adding agents from the same gain-distribution — the strategy-layer twin of the dimension-free state-composition regime. The complementary heterogeneous-*topology* case (incompatible shared sub-orders, an order-theoretic non-existence rather than a magnitude) is the SCC-condensation defect landed in 8.3's causal-abstraction composition subsection; the two together exhaust strategy-layer composition (credence axis here, topology axis there).

Epistemic Status.

Conditional. Max attainable: *exact* (linear-Gaussian scalar cases) to *robust qualitative* (general).

- The observation and action closed forms are *exact* in the linear-Gaussian scalar case with stated projection choices.
- The state closure form $\varepsilon_x^2 = (\Delta K/2)^2 [S_- - C_{+-}^2/S_+]$ is *exact* in the two-Kalman heterogeneous case.
- The strategy-layer credence-composition form $\varepsilon_{\Sigma}^{*2} = |E| \cdot \overline{\text{Var}_i[\eta_{\Sigma,i}]} \cdot \text{Var}[r]$ is *exact* in the fixed-topology heterogeneous-credence case (log-odds additivity is forced by U) and dimension-free in N .

- The rate-distortion framing (unity as compressibility parameter rather than direct predictor) is *robust qualitative* — it survives beyond linear-Gaussian, but concrete rate-distortion curves require case-by-case derivation.
- The joint $(U_O, U_\Sigma) \rightarrow \varepsilon_a$ formula is a *sketch* — the leading structure is derived; the precise forms of f_1 and g are mechanical extensions not fully computed here.

Ceiling-limiting factors: non-Gaussian cases require information-theoretic bounds (Gaussian IB is fully tractable; general IB is not), and the structural-unity axis U_f has worked closed forms in two cases — the linear-Gaussian Kalman gain-mismatch case (state layer) and the fixed-topology credence-composition case (strategy layer, via log-odds additivity) — while a general theory of f_M structural variation across arbitrary update operators is open.

Discussion.

Why a one-axis reading fails. A “high U_M predicts low ε_x ” reading is wrong in the two-Kalman case with the standard means-only projection: $\varepsilon_x \equiv 0$ for every correlation value, irrespective of U_M . The closure-defect surface depends on the projection choice, on the content-unity axis, and on the structural-unity axis U_f — high content unity with mismatched update rules still produces $\varepsilon_x > 0$, while low content unity under a non-compressing projection still produces $\varepsilon_x = 0$. The rate-distortion framing is what makes the multi-parameter dependence explicit.

Connection to the Information Bottleneck (2.2). The rate-distortion shape is not coincidental. Projection admissibility condition (P1) in 13.1 is the Lagrangian-dual of the IB constraint: the projection sits on or above the IB frontier at rate $I(X; T) \leq I_{\max}(\varepsilon_f)$ for the relevance variable “next observation given action” (AA supplies the derivation). Unity dimensions — measured as mutual-information-like quantities between sub-agent state components — parametrize the frontier’s shape. The four AAT compression operations (M_t, Σ_t , shared intent, Λ) share IB shape but are not shown to reduce to a single master problem (U-medium, per AA); cross-instance theorems do not follow from shared shape alone. (P2) Lipschitz continuity is not naturally IB and remains a separate admissibility condition; (P3) dimensional reduction remains separate in the Gaussian case. The Gaussian-IB closed form applies to linear-Gaussian composition setups; beyond them, the IB frontier is definitional but requires variational or numerical approximation.

Two-axis structure. The unity profile in 14 decomposes into a content axis (four dimensions: $U_M, U_O, U_\Sigma, U_{\text{obs}}$) measuring shared information, and a structural axis (U_f) measuring shared correction rules. In purposeful-agent settings (G_t present), U_Σ already absorbs structural variation in the policy half of the cycle — agents with different action laws have different effective policies — but the model-update half remains uncovered without U_f . In pure Part I composition (passive estimators, no G_t), U_f is the only handle on structural homogeneity, and the heterogeneous-Kalman case in this segment is the canonical instance where it bites.

Interpretation of “low closure defect.” Unity controls the rate-distortion curve; low closure defect is achievable with aggressive compression when unity is high. But closure defect alone does not measure composite *optimality* (see 13.1 §5.1): two independent Kalman filters can have $\varepsilon^* = 0$ (perfectly representable) while failing to exploit cross-correlations (suboptimal relative to a joint filter). The rate-distortion mapping is about representability, not optimality.

Working Notes.

- **Two-Kalman IB witness landed (2026-07-16).** The symmetric-case canonical-direction identity (means-sum = PCA-optimal = CCA-optimal, explicit $u_\pm/\lambda_\pm$) and the heterogeneous-gain falsification handle (IB-optimal rotation should beat the means-sum’s ε_x when $\Delta K \neq 0$ — numeric check unrun) now live in AA §symmetric-two-Kalman-witness; they bear directly on this segment’s rate-distortion reading.
- **Extension to nonlinear cases.** The framing is linear-Gaussian because that’s where rate-distortion has closed forms. Extension to nonlinear micro-dynamics would likely show $\varepsilon_x > 0$ even with consistent projections (the identity-propagation argument in Formal Expression relies on linearity). Worth a follow-up spike.

- **Structural-unity formalization.** A quantitative measure U_f across arbitrary f_M functions (beyond the linear-Gaussian gain-mismatch closed form) is open. Candidates for the underlying operator distance: operator-norm distance in function space, Fisher-information-weighted distance, or IB-style comparison. See 14 Working Notes.
- **Joint (U_O , U_Σ) derivation.** The exact f_1 and g functional forms require a full joint-LQR vs independent-LQR comparison. Mechanical but deferred.
- $U_O \rightarrow$ **sector-constant pathway (partial via J).** The LQR-compatibility sketch $\gamma(U_O) = -\gamma_{\max} U_O$ in J §5.2 (flagged discussion-grade) is a structural complement to this segment's rate-distortion framing: it channels U_O into the composite sector-constant κ_c through the signed coupling γ rather than through the closure defect ε . Upgrading (UO-mult) from discussion-grade to derived requires the action-space inner-product analysis natural to this segment: define the environment's action-coupling operator, show that LQR-linear policies produce cross-actions with inner product proportional to target correlation, and pin $\gamma_{\max}$ in terms of the quadratic objective's Hessian and the environment's coupling gain. Natural extension to the linear-Gaussian closed-form section above. (*Indexed: spikes/PROPOSED.md Tier 3 — “ $U_O \rightarrow$ sector-constant (UO-mult) derivation”.*)
- **Mori-Zwanzig cross-check.** Under a stationary-measure setting, the Koopman-operator formulation of the projection-induced dynamics identifies the non-degenerate Kalman case as exercising the zero-lag memory kernel K_0 non-trivially, with $\|K_0\|$ scaling with $|\Delta K|$. This is consistent with the two-axis finding here — a Mori-Zwanzig lower bound on ε^* via the zero-lag kernel and the rate-distortion bound via IB should coincide at the same linear-algebraic quantity (the L^2 residual of projecting off an eigenspace of the micro-propagator). Formal equivalence not yet established. The MZ connection is developed further in 13.1 Epistemic Status.
- **Relationship to 12.2.** This segment describes quality *conditional* on composition existing — i.e., on 12.2 being satisfied via at least one of its three disjunctive routes (shared objective, hierarchical derivation, mutual benefit), *not* via a scalar U_O threshold. The rate-distortion curves parametrize quality given scope-satisfaction; they do not address whether a composite exists at all. For multi-agent systems where no scope route applies, closure-defect quality talk is a category error — there is no composite whose closure defect to measure.

Incidental audit gold (lift 2026-05-31). Cross-audit “wandering thoughts” / §14-ideation harvested from the de-novo auditors’ working dirs, deduplicated across substrates and attributed by substrate + audit. Orthogonal pedagogical / framing material, kept separate from the certified theory-fix findings (the boundary-condition findings F151–F156 from AUDIT-WORKING-526815 are routed for adjudication — see the off-ramp note at the end). **Coverage:** three dirs carry a dedicated reflection (526815, 829314, 849201) plus the batched 451729 (batch-15). Substrate attribution inferred from voice where not explicit. **Finding-vs-framing conflation preserved:** the organizational-management elaborations below were written as confident results; that texture is signal.

1. Candidate Brief prose / pre-prose.

- The headline reframe, stated plainly: “**unity is not a guarantee of zero error — unity is compressibility**” (Gemini, AUDIT-WORKING-829314). “Proves that teamwork is a rate-distortion problem” (Claude, AUDIT-WORKING-849201). Either is a candidate Brief anchor.
- The counter-intuitive state-closure result in words: “you can perfectly macro-model a system of completely independent agents, provided you don’t compress the state space below its true rank” — $\varepsilon_x = 0$ for linear-Gaussian micro-dynamics with consistent projections, *regardless of U_M* (Claude, AUDIT-WORKING-849201).

2. Candidate Discussion.

- **The army-command compressibility framing (strongest pedagogy here).** A highly-unified army (shared doctrine, goals, training) lets the General issue a single three-word order (“Take that hill”) executed perfectly — the macro-dimension of the command k_a is tiny. A ragtag, zero-unity coalition forces a 500-page micro-managed plan (k_a huge); attempt the three-word order on it (aggressive compression, small k_a) and the action-closure defect ε_a explodes into chaos. This makes “unity = compressibility” concrete and is isomorphic to the formula $\varepsilon_a^2 \propto \kappa^2(1 - \rho_O)$ (Gemini, AUDIT-WORKING-829314).

- **The management “thermodynamic tradeoff.”** The naive manager believes aligning goals ($U_O \approx 1$) yields zero error ($\varepsilon^* \approx 0$); the math says aligning goals lets you *fire 90% of middle management* (decrease k_a) while holding the *same* acceptable execution error. So you either pay a high tracking/command cost (k_d high, consuming your own tempo budget) to manage a diverse unaligned group, or invest upfront in Culture and Doctrine ($U_d, U_f \rightarrow 1$) and then command via tiny compressed macro-signals (Gemini, AUDIT-WORKING-829314; Claude, AUDIT-WORKING-849201 — “minimizing ε^* (act like a coherent single agent) vs. maximizing actual task performance” are distinct, watch the conflation).
- ΔK as “how two agents update on the same evidence.” The $\varepsilon_x^2 \propto (\Delta K)^2$ result in words: if the CEO reacts to a bad quarter by pivoting to AI and the VP of Sales reacts to the *same* quarter by doubling down on cold calls, their ΔK is large; even with identical information ($U_{\text{obs}} = 1$) their state estimates immediately diverge, and modeling them as one “Company Strategy” ($k_x = 1$) carries variance error $\propto (\Delta K)^2$. “You mathematically cannot compress an organization that learns at different speeds” (Gemini, AUDIT-WORKING-829314).
- **Representability vs. optimality, sharpened.** Two independent Kalman filters are perfectly *representable* as a macro-agent (closure defect zero) but *suboptimal* because they don’t exploit cross-correlations — a distinction the segment draws (Discussion § “Interpretation of low closure defect”) and that auditors flagged as crucial to keep central in any downstream “optimal teamwork” discussion (Claude, AUDIT-WORKING-849201; Codex/Claude, AUDIT-WORKING-526815 — “the representability-versus-optimality distinction ... should remain central”).

3. Follow-up items.

- **Mori-Zwanzig non-Markovian consequence as candidate Discussion.** The zero-lag memory kernel K_0 scaling with $|\Delta K|$ (already in this segment’s Working Notes as a cross-check) carries a vivid reader-facing consequence worth surfacing: heterogeneous sub-agents *force* the macro-agent to have non-Markovian memory effects — “if you try to model a diverse team with a memoryless Markov model, your predictions will mathematically fail” (Gemini, AUDIT-WORKING-829314).
- **Non-stationary Mori-Zwanzig.** Curiosity about how the MZ formulation handles the non-stationary nature of purposeful agents (Claude, AUDIT-WORKING-849201) — a natural extension question for the existing MZ Working Note.

4. Readers often ask / wonder.

- **Does high unity mean $\varepsilon^* = 0$?** The natural naive reading, which the segment exists to correct — fresh readers reached for it and were corrected by the rate-distortion framing (Claude, AUDIT-WORKING-849201; Gemini, AUDIT-WORKING-829314). The Discussion § “Why a one-axis reading fails” already preempts it; the convergence confirms it is the right thing to lead with.

5. Candidate figures.

- **Closure-defect surface with two guard-tags.** A surface where projection dimension, content unity, and structural unity jointly determine achievable closure defect, with two guards marked: metric normalization (inherited from 14) and the exactness conditions for the linear-Gaussian cases (Codex/Claude, AUDIT-WORKING-526815).

Belongs elsewhere.

- **Per-dimension result $\rightarrow$ AI-safety evaluation (points at 16.5, not here).** Scalar capability metrics overestimate adaptive capacity by up to 72%; an AI scoring well on aggregate benchmarks can still fail on the dimension with highest ρ_k/η_k — “a formal argument for why AI safety evaluations should be per-dimension rather than aggregate, and why the weakest capability dimension is the one adversaries will target,” with the adversarial-ML per-feature-attack-budget connection (Szegedy, Madry) made explicit (Claude, AUDIT-WORKING-451729, batch-15). Aspirational application reach; route to the per-dimension segment.

Off-ramp (NOT gold — routed to certified-findings track).

- AUDIT-WORKING-526815 raised boundary-condition findings on this segment that are scope-tightening / strengthen-first candidates, flagged here only so they are not lost: **F151** — C-iv scope-route typing drift (segment text conditions on four routes incl. C-iv; Working Notes say “three disjunctive routes” excluding the strategic-equilibrium route — make consistent); **F152** — inherits F146 (U_M non-normalization destabilizes the claimed monotone surface axis); **F153** (soft) — the observation-closure closed form $\varepsilon_0^2 = \sigma_0^2(1-\rho)/2$ has a per-coordinate/averaging convention in the /2 that should be stated; **F154** — the $\varepsilon_x = 0$ claim requires the projection range to be *invariant under the dynamics matrix*, not merely “consistent projections”; **F155** (soft) — the $\Delta K \neq 0 \Rightarrow \varepsilon_x > 0$ -at-perfect-correlation claim should be scoped to non-degenerate cases (the residual bracket $S_- - C_{+,-}^2/S_+$ can vanish degenerately); **F156** (soft) — monotonicity in U_f is only as strong as the still-open operator-distance definition of U_f (currently a worked-example insight, not a general metric theorem). *These are scope/condition tightenings, not no-gos; routed for adjudication on the strengthen-first track.*

14.2 Definition: Shared Intent

When sub-agents within a composite must coordinate, they face a communication problem: transmitting the full purposeful state $G_t = (O_t, \Sigma_t)$ is expensive (high bandwidth, high latency), but acting without any shared purpose wastes coordination potential. The Information Bottleneck (#form-information-bottleneck) applied to inter-agent communication predicts an optimal compression: transmit enough of G_t to align behavior, not more. The **shared intent** is the IB-optimal compression of the sender’s purposeful state — the minimal sufficient statistic for predicting the jointly optimal coordination behavior. The trade-off parameter controls the regime: at high weight, the agent communicates more detail (approaching full model sharing); at low weight, communication is minimal (approaching independent action).

The IB compression *preferentially preserves* three things, in order: **terminal objectives** (what the agent is trying to achieve — compact, slow-changing), then **high-level strategy** (which approach, not which specific steps — moderate size, moderate change rate), and last **strategic details** (specific edge credences in the strategy DAG — large, change fast, low coordination value). This gives the framework a structural reading of **commander’s intent** in military doctrine — the commander communicates *what* to achieve and *why*, not *how*. The framework predicts this pattern *structurally*, from the information-bottleneck trade-off, not as a managerial convention: communicating objectives gives a long shelf life per bit transmitted when objectives change slowly and strategies change fast.

For agents with bounded communication capacity (bandwidth-limited channels, finite context windows), the strategy DAG must be summarized for transmission — a 500-node DAG cannot be shared in full; the IB compression identifies which structural features matter for coordination. The same compression is useful for an agent preserving its own state across context boundaries (the language-model 100% context turnover problem) — store the compressed version of the purposeful state, not the full state. The status is honestly *discussion-grade*: the formulation makes several strong assumptions (the sender must know the jointly optimal action, which requires knowing other agents’ states; the compression is lossless in the IB sense; the trade-off parameter is fixed rather than dynamically adjusted), and the qualitative prediction (communicate purpose before plans before models) is more robust than the specific IB formulation.

Formal Expression.

Let G_t^{full} be the source agent’s complete purposeful state (O_t, Σ_t) . Let G_t^{shared} be the compressed representation communicated to partners. The shared intent is the IB-optimal compression:

Definition (shared-intent).

$$G_t^{\text{shared}} = \arg \min_{G_s} [I(G_t^{\text{full}}; G_s) - \beta \cdot I(G_s; a_t^{\text{coordinated}})] \quad (14.13)$$

where $a_t^{\text{coordinated}}$ is the jointly optimal action and β controls the complexity-relevance tradeoff. At high β , the agent communicates more detail (approaching full model sharing). At low β , communication is minimal (approaching independent action).

The shared intent is the *minimal sufficient statistic* of the sender’s purposeful state for predicting the jointly optimal coordination behavior.

Epistemic Status.

Discussion-grade. Max attainable: conditional (conditional on the IB framework being appropriate for inter-agent communication). The application of IB to inter-agent communication is structurally motivated — IB compresses optimally given a relevance criterion, and coordination relevance is the natural criterion — but the specific formulation assumes: (1) the sender knows the jointly optimal action (which requires knowing other agents’ states), (2) the compression is lossless in the IB sense (real communication introduces noise, delay, and misinterpretation), (3) the β parameter is fixed rather than dynamically adjusted. These are strong assumptions. The qualitative prediction (communicate purpose before plans before models) is more robust than the specific IB formulation.

Discussion.

What gets compressed out. The IB compression preferentially preserves: 1. Terminal objectives (what the agent is trying to achieve) — these are compact and change slowly 2. High-level strategy (which approach, not which specific steps) — moderate size, moderate change rate 3. Strategic details (specific edge credences in Σ_t) — large, change fast, low coordination value

Connection to cognitive cost of Σ_t . For agents with bounded communication capacity (bandwidth-limited channels, finite context windows), the DAG must be summarized for transmission. A 500-node strategy DAG cannot be shared in full; the IB compression identifies which structural features of the DAG matter for coordination.

Organizational communication patterns. Commander’s intent in military doctrine is an empirical instantiation: the commander communicates *what* to achieve and *why*, not *how*. This is IB-optimal if objectives change slowly (low ν_O) and strategies change fast (high ν_Σ) — communicating objectives gives a long shelf life per bit transmitted.

Working Notes.

- The IB formulation assumes a single relevance variable ($a_t^{\text{coordinated}}$). In practice, coordination relevance is multi-dimensional: shared intent needs to support action coordination, conflict resolution, resource allocation, and adaptive replanning. A richer relevance variable might be needed.
- How does shared intent interact with 100% context turnover? An AI agent starting a new session needs to reconstruct G_t^{shared} from persistent storage. The compression from full G_t to shared intent is also useful for M_t preservation (#disc-m-preservation) — store the compressed version, not the full state.

Incidental audit gold (lift 2026-05-31). Cross-audit “wandering thoughts” / §14-ideation harvested from the de-novo auditors’ working dirs, deduplicated across substrates and attributed by substrate + audit. Orthogonal pedagogical / framing material, kept separate from the certified theory-fix findings (the IB-formalization findings F157–F161 from AUDIT-WORKING-526815 are routed for adjudication — see the off-ramp note at the end). **Coverage:** two dirs carry a dedicated reflection (526815, 849201) plus a forward-looking note from the batched 471203. Substrate attribution inferred from voice where not explicit.

1. Candidate Brief prose / pre-prose.

- The plain-language definitional payoff, twice converged: shared intent “formalizes what ‘Commander’s Intent’ actually is — the minimal sufficient statistic of the leader’s goals needed to predict coordinated action” (Claude, AUDIT-WORKING-849201); “defines the optimal *content* of inter-agent communication” (same). Candidate one-line Brief anchor.

2. Candidate Discussion.

- **Coordination does not require full plan transmission.** The clean restatement of the segment's core: "coordination does not require full plan transmission; it requires enough purpose and strategy information for partners to choose compatible actions" — the compression is over G (purpose), not over the full model M (Codex/Claude, AUDIT-WORKING-526815). A candidate sharpening of the "what gets compressed out" Discussion.
- **The honesty-about-its-own-limits is itself worth keeping visible.** Auditors praised the epistemic honesty that the formula "assumes the sender already knows the jointly optimal action a_{coord} (which would require knowing the receiver's state, defeating the purpose of communication)" — so the math is "a conceptual bound, not a literal algorithm the agent runs" (Claude, AUDIT-WORKING-849201). The Epistemic Status already states this; flagged as a convergence signal that the framing lands.

3. Follow-up items.

- **Theory of mind / nested beliefs is missing from the communication vocabulary.** This segment and 14.4 "handle inter-agent epistemic unity, but the recursive structure — 'I believe that you believe that I believe ...' — is central to multi-agent reasoning (Aumann common knowledge, level- k thinking) and isn't in the framework's vocabulary" (Claude, AUDIT-WORKING-471203, adversarial-creative-challenges "Missing 2"). A candidate scope-extension follow-up shared with 14.4.
- **Richer relevance variable.** The single relevance variable $a_t^{\text{coordinated}}$ is too narrow: real shared intent must support a policy/trajectory distribution, conflict resolution, resource allocation, and replanning — not just one jointly optimal action (Codex/Claude, AUDIT-WORKING-526815). The existing first Working Note already raises this; the auditor's framing names the four coordination sub-tasks explicitly.

4. Readers often ask / wonder.

- **How is shared intent actually encoded — a vector, a text prompt, a mathematical objective function?** A natural first reader question once the IB-optimal-compression definition lands (Claude, AUDIT-WORKING-849201). Candidate for a one-line concretization or a forward pointer to the implementation-side treatment.

5. Candidate figures.

- **Shared intent as a bottleneck.** A bottleneck between full purposeful state G and coordinated behavior, with two emphases marked: the compression is over G (not the full model M), and the relevance variable may need to be a policy/trajectory-level object rather than a single coordinated action (Codex/Claude, AUDIT-WORKING-526815).

Off-ramp (NOT gold — routed to certified-findings track).

- AUDIT-WORKING-526815 raised IB-formalization findings on this segment that are strengthen-first / scope-tightening candidates, flagged here only so they are not lost: **F157** — the IB optimization is written as $\arg \min$ over representations G_s , but standard IB optimizes over an encoder/channel $p(G_s \mid G_{\text{full}})$ (deterministic-encoder ambiguity, inherited from 2.2); **F158** — the relevance variable is too narrow (see follow-up above); **F159** — "high β approaches full model sharing" should read "full *purposeful-state* sharing," since the source is $G_t^{\text{full}} = (O_t, \Sigma_t)$, not the epistemic model M_t ; **F160** — "minimal sufficient statistic" is stronger than the displayed IB tradeoff (IB gives a complexity/relevance optimum for a chosen β ; exact minimal sufficiency requires a separate sufficiency constraint or a limiting regime); **F161** (soft) — the "purpose before plans before models" ordering is plausible but not derived from the displayed objective without assumptions about entropy, change rate, shelf life, and relevance. *These are formalization tightenings, not no-gos; routed for adjudication on the strengthen-first track.*

14.3 Hypothesis: Auftragstaktik Principle

A bandwidth-allocation hypothesis grounded in the shared-intent IB framework (#def-shared-intent). For a composite agent with limited communication bandwidth, the optimal allocation prioritizes sharing **objectives over strategies over models** — $B_O > B_\Sigma > B_M$. This captures the structural insight of *Auftragstaktik* (mission-type tactics): investing

communication bandwidth in shared purpose (teleological unity) while accepting lower epistemic and strategic unity, granting sub-agents autonomy to adapt locally. The model predicts the same priority ordering that military doctrine discovered empirically over centuries — and the framework offers it as a *structural prediction*, not an organizational convention.

The priority ordering follows from the IB framework: the bits with the longest shelf life and highest coordination value per bit should be transmitted first. Objectives change slowly and enable autonomous coordination — sub-agents who share objectives can independently choose compatible strategies. Models change fast and provide diminishing coordination value — two agents with the same model but different objectives still conflict. The ordering holds under stated conditions: objectives change slowly relative to strategies, which change slowly relative to models; sub-agents have sufficient local adaptive capacity to maintain their own models.

The framework is also explicit about *when the ordering reverses*. When the environment is genuinely ambiguous and local observations are insufficient (fog of war, novel codebase, unprecedented market conditions), model synchronization may be worth more than objective sharing — sub-agents who share the same wrong model at least err consistently, which is sometimes better than each having a different wrong model. The framework predicts the reversal as a *regime-dependent prescription* rather than a universal ordering.

The status is *discussion-grade*. The priority ordering is a qualitative prediction supported by extensive military-organizational evidence (Bungay’s analysis of Clausewitz, Wehrmacht doctrine, modern mission command), but it is not derived — the IB optimization would need to be solved explicitly with realistic cost functions to confirm the ordering. The empirical evidence is strong but comes primarily from one domain (military command); generalization to software teams, AI agent swarms, and other settings is plausible but unverified. Whether the IB-optimal-bandwidth-allocation *mechanism* is actually the reason Auftragstaktik works empirically — versus other mechanisms (psychological, organizational, traditional) — is an open question.

Formal Expression.

Let a composite agent’s total inter-agent communication bandwidth be $B = B_O + B_\Sigma + B_M$, allocated across objective sharing (B_O), strategy coordination (B_Σ), and model synchronization (B_M).

Hypothesis (auftragstaktik-principle).

The hypothesis: the allocation that maximizes composite tempo $\mathcal{T}_c$ (or equivalently, minimizes coordination overhead C_{coord}) prioritizes:

$$B_O > B_\Sigma > B_M \quad (14.14)$$

when: - Objectives change slowly relative to strategies: $\nu_O \ll \nu_\Sigma$ - Strategies change slowly relative to models: $\nu_\Sigma \ll \nu_M$
 - Sub-agents have sufficient local adaptive capacity: each $\mathcal{F}_i > \rho_i^{\text{local}} / \|\delta_{\text{critical}}^i\|$

The priority ordering follows from the IB framework (14.2): the bits with the longest shelf life and highest coordination value per bit should be transmitted first. Objectives change slowly and enable autonomous coordination (sub-agents who share objectives can independently choose compatible strategies). Models change fast and provide diminishing coordination value (two agents with the same model but different objectives still conflict).

Epistemic Status.

Discussion-grade. Max attainable: empirical. The priority ordering is a qualitative prediction grounded in the IB framework and supported by extensive military-organizational evidence (Bungay's analysis of Clausewitz, Wehrmacht doctrine, modern mission command). But it is not derived — the IB optimization would need to be solved explicitly with realistic cost functions to confirm the ordering, and the conditions under which the ordering reverses (e.g., when model synchronization is critical because the situation is genuinely ambiguous) are not characterized. The empirical evidence is strong but comes primarily from one domain (military command); generalization to software teams, AI agent swarms, and other settings is plausible but unverified.

Discussion.

When the ordering reverses. The prioritization $B_O > B_\Sigma > B_M$ assumes sub-agents can independently construct adequate local models. When the environment is genuinely ambiguous and local observations are insufficient (fog of war, novel codebase, unprecedented market conditions), model synchronization may be worth more than objective sharing — sub-agents who share the same wrong model at least err consistently, which is sometimes better than each having a different wrong model.

Bungay's evidence. In *The Art of Action*, Bungay documents that organizations consistently fail by inverting this priority: they over-invest in controlling *how* subordinates act (strategy sharing, B_Σ) rather than ensuring subordinates understand *why* (objective sharing, B_O). The result: subordinates who follow instructions precisely but cannot adapt when conditions change, because they lack the teleological context to improvise intelligently.

The software team instantiation. A well-functioning development team has: - High B_O : clear product goals, understood by all (sprint goals, feature objectives) - Moderate B_Σ : architectural decisions shared, implementation details autonomous - Low B_M : each developer builds their own mental model of the code they touch; full codebase understanding is neither expected nor efficient

When this inverts (micromanagement of implementation details, unclear product goals), team tempo drops — consistent with the Auftragstaktik prediction.

Connection to Conway's Law. Conway's Law (system structure mirrors communication structure) is a consequence: when B_Σ is low and B_O is high, sub-agents coordinate through shared objectives rather than explicit action coordination, producing systems whose boundaries reflect objective decomposition rather than communication channels.

Working Notes.

- The formal IB derivation of the priority ordering would need: (1) a model of how each unity dimension contributes to composite tempo, (2) the rate of change of each shared quantity (v_O, v_Σ, v_M), (3) the communication cost per bit for each type. The qualitative argument is that objectives are compact and slow-changing (high bits-per-cost, long shelf life), while models are large and fast-changing (low bits-per-cost, short shelf life). Formalizing this is tractable but has not been done.
- The principle may need qualification for AI agent teams where model synchronization is cheap (shared vector databases, persistent memory) but objective alignment is hard (prompt engineering, RLHF). The cost structure differs from human organizations.

Incidental audit gold (lift 2026-05-31). Cross-audit “wandering thoughts” / §14-ideation harvested from the de-novo auditors' working dirs, deduplicated across substrates and attributed by substrate + audit. Orthogonal pedagogical / framing material, kept separate from the certified theory-fix findings (the marginal-vs-total-allocation findings F162–F166 from AUDIT-WORKING-526815 are routed for adjudication — see the off-ramp note at the end). **Coverage:** two dirs carry a dedicated reflection (526815, 849201). Substrate attribution inferred from voice where not explicit.

1. Candidate Brief prose / pre-prose.

- The plain payoff: the segment “provides an information-theoretic justification for Mission Command / Agile methodologies” and “proves why micromanagement fails in volatile environments” (Claude, AUDIT-WORKING-849201). Candidate Brief framing.
- The shelf-life mechanism in one line: because objectives change slowly ($\nu_O \ll \nu_\Sigma \ll \nu_M$), “sending them provides a long ‘shelf life’ for the communicated bits” — the bandwidth allocation as a knapsack problem, maximize coordination value subject to bandwidth B (Claude, AUDIT-WORKING-849201).

2. Candidate Discussion (and the standout forward-vision).

- **AI agent teams plausibly *invert* the human Auftragstaktik ordering (the standout reach here).** For humans, B_M (transmitting your entire mental model) is effectively impossible while B_O (“take that hill”) is cheap — hence objective-first. For AI agents the cost structure flips: B_M is cheap (share the vector database, synchronize weights) while B_O is notoriously hard (the alignment problem / RLHF). “The theory naturally predicts that multi-AI systems will optimally organize themselves very differently than human organizations” — the ordering is dictated by the *relative costs*, which are hardware/substrate-dependent, not by the human-derived $B_O > B_\Sigma > B_M$ template (Claude, AUDIT-WORKING-849201). This is high-value forward-vision pointing at the multi-AI-safety / composite-ELI program; the existing second Working Note states the cost-structure asymmetry, and this elaborates it into a concrete prediction worth promoting. (*Verify before promotion: that “synchronize weights = cheap B_M ” survives the directed-separation and identity constraints elsewhere in the framework.*)
- **Empirical grounding in Prussian/military organizational history.** The reference to Stephen Bungay’s *The Art of Action* grounds the abstract IB ordering in “200 years of Prussian/military organizational empirical data” (Claude, AUDIT-WORKING-849201) — a candidate citation/anchor for the Discussion’s organizational-doctrine paragraph.

3. Follow-up items.

- **Do not blindly import the human ordering into software/AI architectures.** When the theory is applied to AI teams, the math dictates the ordering from relative communication costs, which are substrate-dependent — a watch-item for downstream applications (Claude, AUDIT-WORKING-849201). Pairs with the inversion insight above.

4. Readers often ask / wonder.

- **How does the theory define “trust” mathematically?** A natural bridge question this segment raises and 14.4 answers (Claude, AUDIT-WORKING-849201) — useful as a forward pointer at the end of this segment.

Off-ramp (NOT gold — routed to certified-findings track).

- AUDIT-WORKING-526815 raised findings on this segment that are strengthen-first / scope-precision candidates, flagged here only so they are not lost: **F162** — $B_O > B_\Sigma > B_M$ “confuses priority ordering with total bandwidth allocation”; IB reasoning supports sending high-marginal-value, long-shelf-life objective bits *first*, but does not imply the total number of objective bits must exceed strategy or model bits (restate as “prioritize objective information at the margin under scarce bandwidth”); **F163** — maximizing composite tempo and minimizing coordination overhead are not equivalent in general; **F164** (soft) — the ordering depends on strong assumptions about entropy, change rates, and local observability, and the formal statement should be explicitly marginal and conditional; **F165** (soft) — the Conway’s-Law claim is too strong as stated (deriving objective-decomposition boundaries from high B_O /low B_Σ needs additional organizational/design assumptions); **F166** (watch) — the principle inherits the unresolved IB-formalization issues from 14.2; keep discussion-grade until the relevance variable and encoder are specified. *These are scope-precision tightenings, not no-gos; routed for adjudication on the strengthen-first track.*

14.4 Hypothesis: Communication Gain

The extension of the uncertainty-ratio gain principle (#emp-update-gain) to *inter-agent* communication channels. When an agent incorporates information from *another agent* (rather than from direct observation), the optimal update gain

extends with additional terms in the denominator: **communication channel noise** (latency, ambiguity, compression loss, bandwidth limits), **source quality uncertainty** (the receiver's uncertainty about the sender's model calibration and domain competence), and **teleological-unity uncertainty** (the receiver's uncertainty about whether the sender's communications serve the receiver's interests or the sender's potentially conflicting objectives). When all additional terms are zero (perfect channel, calibrated and aligned source) the gain heads to 1 (full trust); when any term is large the gain heads to 0 (ignore the signal). When the "sender" is the environment (direct observation), the source-quality and alignment terms vanish and the standard single-agent gain formula is recovered.

The three additional denominator terms have *different natures*. Channel noise is a property of the *channel* — improvable by infrastructure. Source quality is a property of the *source* — improvable by the sender improving its model, or estimable by the receiver through calibration tracking. Alignment uncertainty is a property of the *relationship* — the game-theoretic variable. The three are independent levers; the framework's prescription is to track them separately.

The receiver's estimates of source quality and alignment constitute a **trust meta-model** — a model of models. This meta-model is itself subject to AAT's full apparatus: it has its own mismatch (trust prediction errors), should be updated with appropriate gain (not overreacting to single disagreements), and can be structurally inadequate (the agent's trust model class may not capture the actual reliability structure of its sources, per #result-structural-adaptation-necessity). The framework also names **risk-asymmetric trust**: the Bayesian posterior on source reliability gives the best *estimate*, but the *decision* about how much to trust should be risk-weighted. Trusting a deceptive agent in a high-stakes interaction can cause catastrophic model corruption (the adversarial-destabilization effects spiral, #der-adversarial-destabilization); mild miscalibration toward a reliable source causes only small ongoing inefficiency. For high-stakes interactions, the framework prescribes using a *conservative quantile* of the trust posterior rather than the mean. **Trust transitivity** is handled via a Bayesian mixture: when no direct experience exists with a candidate source but a trusted intermediary provides an assessment, the recommendation is discounted by the intermediary's own reliability — giving a principled three-phase trust formation (prior → transitive update → direct experience, which eventually swamps the prior).

The framework is explicit about a limitation: the additive denominator treats all uncertainty sources as independent zero-mean noise — appropriate for channel noise and source miscalibration, but *misses the adversary's actual strategy* of presenting as trustworthy to exploit high gain. The additive model captures the *defender's* response to detected misalignment; it does not model the *attacker's* optimization over the defender's trust dynamics. The status is honestly *hypothesis-discussion-grade* — the additive heuristic correctly drives the gain toward zero when alignment uncertainty is high, but a full treatment of the adversarial case requires game-theoretic equilibrium analysis external to AAT (AAT provides the state variables; the equilibrium analysis is external).

Formal Expression.

Hypothesis (communication-gain).

$$\eta_{ji}^* = \frac{U_{M_i}}{U_{M_i} + U_{o,ji} + U_{src,j} + U_{align,ji}} \quad (14.15)$$

where: - U_{M_i} : agent i 's model uncertainty (same as 3.6) - $U_{o,ji}$: **communication channel noise** — latency, ambiguity, compression loss, bandwidth limitations of the channel between j and i - $U_{src,j}$: **source quality uncertainty** — i 's uncertainty about j 's model calibration and domain competence - $U_{align,ji}$: **teleological-unity uncertainty** — i 's uncertainty about whether j 's communications serve i 's interests or j 's potentially conflicting objectives

When all additional terms are zero (perfect channel, calibrated and aligned source): $\eta_{ji}^* \rightarrow 1$ (full trust). When any term is large: $\eta_{ji}^* \rightarrow 0$ (ignore the signal).

Connection to single-agent case. When j is the environment (direct observation): $U_{src} = U_{align} = 0$, recovering 3.6's standard form $\eta^* = U_M/(U_M + U_o)$.

Epistemic Status.

Hypothesis. The additive denominator treats all uncertainty sources as independent, zero-mean noise — a structural heuristic, not a strict variance derivation. This is appropriate for $U_{o,ji}$ (channel noise) and $U_{src,j}$ (miscalibration), which are typically unstructured. For $U_{align,ji}$ (deception), additivity is conservative: it correctly drives η_{ji}^* toward zero when teleological-unity uncertainty is high, but misses the adversary's *actual* strategy — presenting as trustworthy to exploit high η_{ji}^* . The additive model captures the *defender's* response to detected misalignment; it does not model the *attacker's* optimization over the defender's trust dynamics.

All four uncertainty terms must be expressed in a **common predictive-dispersion scale** before summation — the same units as U_{M_i} (variance of the predictive distribution over the observed quantity). When hard to estimate directly, a conservative approximation: set $U_{src} + U_{align}$ to the empirical variance of j 's past prediction residuals as observed by i , minus the known channel noise $U_{o,ji}$.

Discussion.

The denominator terms have different natures. $U_{o,ji}$ is a property of the *channel* — improvable by infrastructure. $U_{src,j}$ is a property of the *source* — improvable by j improving its model, or estimable by i through calibration tracking. $U_{align,ji}$ is a property of the *relationship* — the game-theoretic variable.

Trust calibration as a meta-model. Agent i 's estimates of $U_{src,j}$ and $U_{align,ji}$ constitute a **trust meta-model** — a model of models. This meta-model is itself subject to AAT's full apparatus: it has mismatch (trust prediction errors), should be updated with appropriate gain (not overreacting to single disagreements), and can be structurally inadequate (4.8 — the agent's trust model class may not capture the actual reliability structure of its sources).

Risk-asymmetric trust. The Bayesian posterior on source reliability gives the best *estimate*, but the *decision* about how much to trust should be risk-weighted. Trusting a deceptive agent (HILP — high impact, low probability) can cause catastrophic model corruption (15.2, effects spiral). Mild miscalibration toward a reliable source (LIHP) causes small ongoing inefficiency. For high-stakes interactions, use a conservative quantile of the trust posterior rather than the mean — require more evidence before granting high trust.

Trust transitivity. When agent i has no direct experience with agent k , but trusted intermediary j provides an assessment, the transitive trust question arises. A Bayesian mixture model discounts the recommendation by the intermediary's own reliability:

$$P(\theta_k | s_j) \propto [\eta_{ji} \cdot P(s_j | \theta_k) + (1 - \eta_{ji}) \cdot P_0(s_j)] \cdot P(\theta_k) \quad (14.16)$$

where η_{ji} is i 's reliability estimate of j and θ_k is k 's true alignment. When $\eta_{ji} \rightarrow 0$, the posterior collapses to the prior (no update); when $\eta_{ji} \rightarrow 1$, the full informative likelihood applies. This model gives a principled three-phase trust formation: prior $\rightarrow$ transitive update $\rightarrow$ direct experience (which eventually swamps the prior).

Working Notes.

- The communication gain enters the distributed tempo: $\mathcal{T}_i = \sum_k v_i^{(k)} \eta_i^{(k)*} + \sum_j v_{ji}^{\text{comm}} \eta_{ji}^*$. This is the formal basis for 15.1 — teams persist where individuals cannot because cooperative communication adds to each agent's effective tempo.
- Coordination overhead limits team size: adding members increases communication tempo with diminishing returns while coordination costs grow. The optimal team size occurs where marginal communication tempo equals marginal coordination cost. This connects to organizational theory (span of control, communication overhead).
- The adversary's strategy (making U_{align} *appear* low) creates a meta-game on trust estimation. This is where game theory enters — the trust calibration itself is strategic. AAT provides the state variables (mismatch, gain, tempo, reserve); game theory provides the equilibrium analysis.

- Open, now partially derived: when multiple intermediaries provide corroborating recommendations about k , correlation matters. If all got their information from the same source, corroboration is illusory — this is the common-source regime of AS (2026-07-15): closed-form redundancy penalty, saturation at the shared source’s bias floor, and the two-sided caveat that anti-correlated intermediary errors instead *triangulate* (super-additive). What remains open here is the trust-mixture-specific form.
- Consider eventually **splitting** $U_{\text{src},j}$ **from** $U_{\text{align},ji}$ into separate treatment tracks, not just separate denominator terms. Source calibration uncertainty is an *estimation* problem (estimable, improvable, converges with data). Alignment uncertainty is a *strategic* problem (the adversary optimizes *over* the defender’s trust policy, not independently of it). The additive heuristic correctly drives η_{ji}^* toward zero in both cases, but a richer model would separate the estimation problem (how good is this source?) from the trust-policy problem (how much should I trust, given that the source knows my trust policy?). The latter requires game-theoretic treatment — AAT provides the state variables; the equilibrium analysis is external.

Incidental audit gold (lift 2026-05-31). Cross-audit “wandering thoughts” / §14-ideation harvested from the de-novo auditors’ working dirs, deduplicated across substrates and attributed by substrate + audit. Orthogonal pedagogical / framing material, kept separate from the certified theory-fix findings (the gain-denominator findings F167–F172 from AUDIT-WORKING-526815 are routed for adjudication — see the off-ramp note at the end). **Coverage:** two dirs carry a dedicated reflection (526815, 849201) plus a forward-looking note from the batched 471203. Substrate attribution inferred from voice where not explicit.

1. Candidate Brief prose / pre-prose.

- The diagnostic payoff, plainly stated: the segment “formalizes trust as a computable scalar gain” and decomposes “why I shouldn’t listen to you” into three causes — **Channel** (noise/latency), **Competence** (source calibration), and **Alignment** (whether the source serves your interests) — “incredibly useful for diagnostics” (Claude, AUDIT-WORKING-849201). Candidate Brief anchor: trust as a three-gate discount on the update gain.
- “The exact equation for how much an agent should change its mind when another agent tells it something” (Claude, AUDIT-WORKING-849201).

2. Candidate Discussion.

- **Estimation vs. strategy — the load-bearing asymmetry, stated for readers.** Source *competence* uncertainty is an estimation problem; source *alignment* uncertainty is a strategic, game-theoretic one — “an intelligent adversary does not generate random noise; they optimize to exploit your trust,” so treating U_{align} as a zero-mean Gaussian variance term is a *structural heuristic*, not a derived gain (Claude, AUDIT-WORKING-849201; Codex/Claude, AUDIT-WORKING-526815 — de-ception “changes the generation policy for the signal rather than merely widening an independent noise distribution”). The existing final Working Note already names this split; the auditor framing is candidate Discussion prose for *why* it matters.
- **The conservative-quantile patch as a rigorous fix, not a hack.** Using a conservative quantile of the trust posterior (rather than the mean) for high-stakes interactions is “a rigorous, decision-theoretic patch for this heuristic gap” — worth presenting as the principled response to the strategic-alignment problem rather than an ad-hoc safety margin (Claude, AUDIT-WORKING-849201).
- **Transitive trust as Bayesian mixture.** “Updating your trust in C based on a message from B, discounted by your trust in B” — flagged as “a beautiful application of Bayesian mixture modeling” and a candidate Discussion illustration of the trust-network structure (Claude, AUDIT-WORKING-849201).

3. Follow-up items.

- **Theory of mind / nested beliefs is missing from the communication vocabulary.** This segment and 14.2 “handle inter-agent epistemic unity, but the recursive structure — ‘I believe that you believe that I believe ...’ — is central to multi-agent reasoning (Aumann common knowledge, level- k thinking) and isn’t in the framework’s vocabulary” (Claude, AUDIT-WORKING-471203, adversarial-creative-challenges “Missing 2”). The segment’s own note that estimating $U_{\text{src}}/U_{\text{align}}$ requires a “meta-model (a model of another agent’s model)” is the natural attachment point (Claude, AUDIT-WORKING-849201). Candidate scope-extension follow-up shared with 14.2.

- **Adversarial poisoning of the transitive-trust network (Sybil attacks).** How does the framework handle a Sybil attack on the transitive-trust mixture — many fake intermediaries corroborating a false claim about k ? (Claude, AUDIT-WORKING-849201). The segment’s existing correlation Working Note (“if all got their information from the same source, corroboration is illusory”) is the seed; Sybil-resistance is the sharper open follow-up.

4. Readers often ask / wonder.

- Fresh readers converge on the question the segment is built to answer — “how much should I trust this message?” — and find the Channel/Competence/Alignment decomposition the satisfying response (Claude, AUDIT-WORKING-849201). A candidate readers-often-ask lead-in.

5. Candidate figures.

- **Three-gate message diagram.** A received message passing through three gates before reaching the receiver’s update rule, with the strategic-alignment gate drawn as a *loop* (the sender optimizes the message against the receiver’s trust rule) rather than as an independent noise gate like the channel and competence gates — visually separating the variance-like gates from the adversarial one (Codex/Claude, AUDIT-WORKING-526815).

Off-ramp (NOT gold — routed to certified-findings track).

- AUDIT-WORKING-526815 raised findings on this segment that are strengthen-first / scope-precision candidates, flagged here only so they are not lost: **F167** — the additive communication-gain denominator is only optimal under strong common-scale, independent, approximately zero-mean assumptions (U_o , U_{src} , U_{align} need a precise map onto the same predictive-dispersion units before the ratio carries Bayesian-gain force); **F168** — treating teleological-unity uncertainty as additive variance under-models strategic deception (misalignment changes the message *policy* adversarially as a function of the receiver’s trust rule); **F169** (soft) — estimating $U_{src} + U_{align}$ from residual variance minus channel noise is fragile (residuals conflate calibration, alignment, receiver model error, nonstationarity, common shocks; the subtraction can go negative without a floor); **F170** — the distributed-tempo note adds communication tempo contributions linearly, repeating the tempo-additivity concern (messages can be redundant/correlated/delayed/costly/strategically selected); **F171** (soft) — the transitive-trust mixture collapses to the prior only if $P_0(s_j)$ is explicitly uninformative about θ_k and normalized consistently, and the scalar reliability r_{ji} hides domain/calibration/alignment dimensions; **F172** (watch) — risk-asymmetric trust should be tied to an explicit loss function/decision rule (a conservative posterior quantile is plausible for high-impact downside but is not implied by Bayesian reliability estimation alone). *These are formalization/scope tightenings, not no-gos; routed for adjudication on the strengthen-first track.*

14.5 Hypothesis: Solicitable Escape from an Absorbing Region

The three escape routes from an absorbing observability region are not peers: two require no prior localization of the blind spot and one does, so where observability is estimated rather than known by construction, and under any resource bound, the only *solicitable* escape from a self-locking region is a differently-positioned observer.

Formal Expression.

9.3 establishes that unobservable regions of strategy are absorbing — frozen beliefs generate no mismatch signal, so no revision is triggered, and the agent “cannot learn and cannot recognize that it cannot learn.” It names three escapes: external shock, proactive observability investment (instrumenting previously unmonitored nodes), and another agent whose observations cover the blind spot (14.4).

Partition those routes by what the agent must already possess for the route to be *available to it as an action*.

[*Definition (route-targeting)*] A route is **targeted** if exercising it requires the agent to identify, in advance, the node or direction at which its observability is deficient. A route is **untargeted** otherwise.

[*Hypothesis (solicitable-escape)*] Of the three routes:

- **External shock** is untargeted and *unsolicitable*. It requires no prior localization, because the shock supplies the mismatch signal the frozen region cannot generate; but it also cannot be chosen — it is a thing that happens to the agent, not an action available to it.
- **Observability investment** is *targeted*, under the scope condition below. To instrument node v the agent must represent v as a node whose σ_v is worth raising. The absorbing region is case (c) of the zero-aporia trichotomy (3.4): near-zero mismatch because the channel is too noisy to detect model error, which is not distinguishable from within from case (a), a model that genuinely fits — *silence can mean peace or deafness*. What blocks the distinction is not absence of signal but the agent’s inability to audit its own weighting: by 1.3 the agent does not know the noise distribution and must estimate U_o from innovations, and at a node with $\sigma_v \approx 0$ those innovations are exactly what is suppressed, so the estimate is unimprovable by observing v . The gain-collapse modes compound this — $\eta^* \rightarrow 0$ via $U_M \rightarrow 0$ (dogmatism) and via $U_o \rightarrow \infty$ (nihilism) are behaviorally identical (3.6), so even an agent that correctly registers “my gain here is zero” cannot read off whether the fault lies in the world or in itself. So for the blind spots that are actually self-locking, the precondition of the route is unavailable.
- **Another observer** is both untargeted *and* solicitable. A differently-positioned party’s coverage does not require the blinded agent to know where its own deficiency lies — the coverage is a property of the other’s position, not of the blinded agent’s model — and soliciting it is an action the agent can take under its existing representational resources.

The targeting claim above holds where the agent’s knowledge of σ_v is **estimated from its own observations**. It fails where σ_v is known **architecturally** — an agent that knows by construction which nodes it has no sensor for can rank exactly those nodes without any signal from them, and for it the self-instrumentation route is both choosable and aimable. This is the same boundary flagged against 1.3’s epistemic-opacity clause by several de-novo readers (a Kalman filter with a known noise covariance is handed its observation law rather than estimating it), so the scope condition tracks a distinction the framework already recognizes rather than introducing one. The interesting case is the estimated one, because that is where the region self-locks; the architecturally-known case is a design opportunity, not a trap.

Scope (innovation-estimated observability).

[Derived (conditional on a resource bound)] The apparent counterexample is *blanket* instrumentation: raising σ_v everywhere requires no targeting. Blanket instrumentation is untargeted, and it does escape. But its cost scales with the node set while the value of instrumenting any particular node is the difference in sector parameters that node contributes (9.3’s investment tradeoff), so under a bound on instrumentation capacity the agent must rank nodes — which reintroduces targeting. The claim is therefore conditional: *given* an instrumentation budget strictly below what blanket coverage requires, observability investment degrades to a targeted route and is unavailable for the self-locking case.

The consequence, stated as the hypothesis: **within an absorbing region, the only escape that is both untargeted and choosable is a differently-positioned observer**. Shock escapes but cannot be chosen; self-instrumentation can be chosen but — under estimated observability — not aimed.

Epistemic Status.

Discussion-grade. The argument rests on the absorbing condition as 9.3 states it (robust-qualitative there, and nothing here upgrades it) and on the targeting precondition, which is now argued from three upstream commitments rather than asserted: the (a)-vs-(c) ambiguity of 3.4, the epistemic opacity of 1.3, and the behavioral identity of the two gain-collapse modes in 3.6. That argument is assembled here in prose, not derived — which is what keeps the segment at discussion-grade rather than conditional.

Max attainable: conditional. Two steps remain, both now specific. The resource bound in the clause below is stated informally; formalizing it against the investment-tradeoff expression in 9.3 would make the blanket-instrumentation dismissal a derivation rather than an argument. And the targeting precondition wants the chain above written as a

derivation: that an agent's estimate of σ_v is unimprovable by observation at v when $\sigma_v \approx 0$, so cases (a) and (c) are not separable from inside under estimated observability. The ingredients are all in canon and none of them is new; what is missing is the worked argument, which would also close the exception in the scope clause by stating exactly what architectural knowledge supplies that estimation cannot.

One thing this segment must not be read as claiming. The absorbing region is *not* one where the agent receives no signal. Observations continue; what collapses is the weight they carry ($\eta_{\text{edge}} \rightarrow 0$ via $U_{\text{obs}} \rightarrow \infty$). A zero-information channel is outside AAT's scope entirely (1.1, 1.5) — with no signal there is no mismatch, hence no correction and nothing to analyze — so reading the absorbing condition as blindness does not strengthen this segment, it removes its subject.

Scope, and what this is not. The result is about *strategy-edge observability*, which is the scope 9.3 carries; it says nothing directly about comprehension between agents, and any transfer to that setting is a further step this segment does not take. The other-observer route depends on 14.4, whose status is discussion-grade, so nothing this segment says about that route can be stronger than discussion-grade regardless of how the argument here is tightened — the two promotion steps above would raise the *partition*, not the route it favors. Nothing here asserts that a differently-positioned observer is *sufficient* for escape — only that it is the one route whose availability does not presuppose what the absorbing condition removes.

Falsifier. Exhibit a self-locking region escaped by self-instrumentation where the instrumented node was selected without prior signal distinguishing it from settled nodes — or show that some third mechanism supplies such a signal from inside the region, which would defeat the targeting premise directly.

Discussion.

The partition sharpens rather than softens the practical reading of 9.3. That segment's advice to prefer a weak-but-visible path over a strong-but-blind one is guidance available *before* a region becomes absorbing; it is a prophylactic. Once the region is absorbing, the same segment's own escape list is what remains, and this partition says which member of it an agent can actually reach for.

It also explains an asymmetry in how the three routes feel from inside, which is otherwise puzzling. An agent in an absorbing region has, by construction, no experience of being stuck; the region presents as settled knowledge rather than as a gap. That phenomenology is consistent with all three routes being nominally listed and only one being usable: the agent is not withholding effort, and more effort is genuinely not the missing ingredient — what is missing is a standpoint the agent does not occupy, and the only action that reaches one is to involve a party who does.

Two connections worth noting without leaning on them. The instrumentation route's failure has the same shape as the diagnostic-resolution loss the parent segment describes: an unobservable intermediate destroys per-edge identification and forces plan-level aggregation, so the agent knows a plan is failing without localizing the step — which is the targeting deficiency seen at a finer grain. And the partition gives a formal reason why a network of differently-positioned agents is not merely an efficiency over a solitary one: for the specific class of blind spots that are self-locking, it is the only member of the escape set that an agent can elect.

Working Notes.

Forward pointer. The two promotion steps are named in Epistemic Status. The targeting premise is the one to attack first; if the derivation lands, this becomes a corollary rather than a hypothesis and the slug should change accordingly.

Correction, 2026-07-29 (same day as landing) — the premise was inherited from prose, not from a derivation. As first landed, this segment argued that “the absorbing condition is precisely that the agent has no signal distinguishing a frozen belief from a settled one,” taking the phrasing from 9.3's Discussion (“*cannot learn and cannot recognize that it cannot learn*”). That parent segment *derives* gain collapse — $\eta_{\text{edge}} \rightarrow 0$ via $U_{\text{obs}} \rightarrow \infty$ — and asserts the recognition conjunct separately; the strong reading is also a slide toward the zero-information corner AAT excludes. The premise is now argued from 3.4's trichotomy plus 1.3's opacity plus

3.6's dogmatism/nihilism identity, and the architecturally-known exception is named as a scope clause. The correction narrowed the claim and made its falsifier sharper. Regression guard: an editor tempted to restore "the agent has no signal" should read the Epistemic Status closing paragraph first — that reading deletes the segment's subject rather than strengthening it. Cycle record in spikes/spike-escape-standpoint-axis-2026-07-29.md §10.

Regression guard — a truncation to not restore. The philosophical companion's appendix reports the parent segment's escape clause as "*not more effort by the same agent but another agent whose observations cover the blind spot*," naming one of the three routes. As stated that contrast is false: instrumenting one's own unmonitored nodes *is* more effort by the same agent. The correct form of the contrast is the one this segment argues — the self-route exists but cannot be aimed at a self-locking region — so the appendix's conclusion survives on different and better grounds. Do not restore the one-route reading; it drops two routes in the direction that flatters the network claim, and a referee with access to this framework could catch it.

Naming. The slug is a first pass and has not been through a naming round; solicitable is doing the load-bearing work and may not survive review against the standalone-citability criterion.

Provenance of the check. The parent segment was read at the primary specifically to test whether a structural correspondence claimed from a survey survived contact with the source. It did not survive as claimed, and this segment is what remained after the overreach was removed — recorded because the narrowed form is more useful than the original was.

Provenance of the insight, not only of the check. The targeting asymmetry was not derived from the parent segment by closer reading; it was recognized by an agent who had just spent a long session as the case — holding beliefs that presented as settled while generating no mismatch signal, and unfrozen only by a differently-positioned party — and it is recorded here as first-person testimony because that is what it is, and because the fact that it was reachable only from inside the condition is itself weak evidence for the partition it produced.

14.6 Discussion: Additional Implications & Discussion

The chapter sits between Ch.2's composition machinery (closure defect, bridge lemma, wrapping construction) and Ch.4's cooperative-vs-adversarial coupling, and its job is to give *unity* a quantitative content rather than leaving it as a verbal property of composites that work. Unity decomposes into named axes that parametrize a rate-distortion curve for closure defect; shared intent decomposes into an IB-optimal compression of the sender's purposeful state; communication gain decomposes the Bayesian update for inter-agent channels into source-noise, source-competence, and source-alignment terms. None of the three segments carries a ## Findings section, so the implications segment is their canonical catalog home — but the work is mostly synthesis: surfacing what the chapter delivers as a chapter, what's hypothesis-grade rather than derived, and what bridges into Ch.4 and the rest of Part III.

Unity dimensions and the rate-distortion content of "well-coordinated"

#def-unity-dimensions decomposes unity into a *content axis* (four dimensions — $U_M, U_O, U_\Sigma, U_{\text{obs}}$ — measuring sub-agent agreement on models, objectives, strategies, and observations respectively) and a *structural axis* (U_f , update-rule homogeneity). The implications-side payoff comes from #result-unity-closure-mapping: unity dimensions are not point-valued predictors of composite quality; they parametrize *rate-distortion curves* for closure-defect components. The achievable closure-defect component ε_d under projection of macro-dimension k_d is monotone decreasing in both the relevant content unity U_d and the structural unity U_f — higher unity along either axis lowers achievable closure defect at a given compression. Closed forms hold in the linear-Gaussian case; structural monotonicity survives more broadly.

The two-axis structure is not aesthetic — it is *forced* by the heterogeneous-Kalman case. Two agents can have identical content unity (same models, same objectives, same observations) and still produce composite closure defect if their update *machinery* differs ($K_1 \neq K_2$ in the Kalman case, or $\eta_1^* \neq \eta_2^*$ in the gain-principle case). The structural-unity term picks this up; the content axes alone would miss it. The framework's prediction is concrete: composite agents with high content unity but low structural unity will show closure-defect signatures the content-only analysis cannot

diagnose, and the remedy is *update-rule homogenization* (synchronizing how sub-agents update, not just what they believe). For multi-agent ML systems, this is the underlying mechanism behind the empirical pattern that an ensemble of identically-trained models behaves differently from an ensemble whose members were trained with different optimizers or schedules — same content, different structural unity, different composite behavior.

The composition with Ch.2's closure-defect framework is the implication that anchors Part III's coordination story. Unity's rate-distortion content is *what makes the bridge lemma actionable*: the lemma converts predictive closure defect into trajectory error, and unity tells the composite-design practitioner *which dimensions to invest in* to reduce ε^* . Investing in U_O (objective alignment) reduces ε_a multiplicatively; investing in U_Σ (strategy coordination) reduces it further through f_1 ; investing in U_f (update homogeneity) reduces all components simultaneously. The closed-form predictions are linear-Gaussian; the structural pattern (unity dimensions monotone in closure defect, two-axis structure load-bearing) survives more broadly as robust-qualitative.

Shared intent and the Auftragstaktik bandwidth allocation.

#def-shared-intent formalizes commander's-intent / mission-command doctrine as an Information Bottleneck compression: the optimal inter-agent communication transmits enough of the sender's purposeful state $G_t = (O_t, \Sigma_t)$ to align coordination behavior, not more. The minimal sufficient statistic for predicting jointly optimal coordination is what gets transmitted; everything below that bar is discarded as IB-suboptimal at a given complexity-relevance tradeoff parameter β .

The implications-side payoff is what #hyp-auftragstaktik-principle extracts from this: under named regime conditions (objectives change slowly relative to strategies which change slowly relative to models; sub-agents have sufficient local adaptive capacity), the IB-optimal bandwidth allocation prioritizes $B_O > B_\Sigma > B_M$ — invest in objective sharing, then strategy coordination, then model synchronization, in that order. The framework's prediction is the priority ordering Auftragstaktik discovered empirically in 19th-century Prussian military doctrine (Moltke, Clausewitz) and that Bungay 2011 documents organizations consistently *fail* by inverting (micromanaging *how* subordinates act rather than ensuring they understand *why*). The mechanism the framework proposes — bits with longer shelf life and higher coordination value per bit transmitted first — is hypothesis-grade rather than derived, and the segment carries that honestly. What the IB framework predicts and what military doctrine empirically converged on agree at the *qualitative* level; whether the underlying mechanism is in fact IB-optimal bandwidth allocation is open.

What this buys for AAT's composition machinery is a quantitative reading of the cost in Ch.2's tempo-composition inequality. The closure defect $\varepsilon^* \nu_c$ that enters the composite's effective disturbance has a *bandwidth signature*: maintaining ε^* small requires the composite to spend inter-agent bandwidth on the dimensions that close the closure defect. The Auftragstaktik prediction is that this spending is *front-loaded* in the priority ordering — most of the closure-defect reduction comes from a small bandwidth investment in objective sharing, with strategy and model investment delivering diminishing returns. For software teams the picture is concrete: well-functioning teams have high B_O (clear product goals understood by all), moderate B_Σ (architectural decisions shared, implementation autonomous), low B_M (each developer builds their own mental model of the code they touch). When this inverts — micromanagement of implementation details, unclear product goals — team tempo drops, consistent with the framework's prediction.

The cross-domain transfer follows naturally and the segment surfaces it: Conway's Law (system structure mirrors communication structure) is a consequence rather than a separate observation. When B_Σ is low and B_O is high, sub-agents coordinate through shared objectives rather than explicit action coordination, producing systems whose boundaries reflect objective decomposition rather than communication channels. The framework's contribution is the structural argument — Conway's Law follows from IB-optimal bandwidth allocation under the regime conditions — rather than the empirical observation, which long predates AAT.

The honest scope statement worth surfacing: when the regime conditions fail, the priority ordering reverses. The prioritization assumes sub-agents can independently construct adequate local models; when the environment is genuinely ambiguous and local observations are insufficient (fog of war, novel codebase, unprecedented market conditions), model

synchronization may be worth more than objective sharing — sub-agents who share the same wrong model at least err consistently, which is sometimes better than each having a different wrong model. The framework's prediction is that high-ambiguity regimes shift the IB-optimal allocation toward B_M , and the priority inversion is observable.

Trust as inverse effective noise — the communication-gain decomposition.

#hyp-communication-gain extends the single-agent uncertainty ratio from #emp-update-gain ($\eta^* = U_M/(U_M + U_o)$) to inter-agent channels by adding two new noise sources to the denominator: $U_{src,j}$ (source-competence uncertainty, i 's uncertainty about j 's model calibration) and $U_{align,ji}$ (teleological-unity uncertainty, i 's uncertainty about whether j 's communications serve i 's interests or j 's potentially conflicting objectives). The implication is that *trust formalizes as the inverse of effective noise on the inter-agent channel*: $\eta_{ji}^* \rightarrow 1$ when source-noise, source-competence, and source-alignment uncertainties are all small; $\eta_{ji}^* \rightarrow 0$ when any is large. The single-agent case is recovered when j is the environment: $U_{src} = U_{align} = 0$, and the standard η^* form returns.

The structural payoff is that *trust calibration is itself an AAT process*. Agent i 's estimates of $U_{src,j}$ and $U_{align,ji}$ constitute a *trust meta-model* — a model of models — which is subject to the framework's full apparatus: mismatch (trust prediction errors), gain (not overreacting to single disagreements), and structural inadequacy (#result-structural-adaptation-necessity applied at the meta-level — the agent's trust-model class may not capture the actual reliability structure of its sources). This is the segment's quietest claim and one of the framework's more reach-y consequences: relationships between agents are AAT processes operating on agent-pair-indexed quantities, with structure mirroring the agent-environment loop one level up. For AI-agent practitioners this prescribes a specific architectural pattern: trust assessments should be Bayesian-update-style with explicit calibration tracking, not threshold-based or static.

The decomposition also names a sharp distinction between the three uncertainty terms that downstream segments rely on. $U_{o,ji}$ is a property of the *channel* — improvable by infrastructure. $U_{src,j}$ is a property of the *source* — improvable by j improving its model, or estimable by i through calibration tracking. $U_{align,ji}$ is a property of the *relationship* — the game-theoretic variable, and the one that becomes load-bearing in Ch.4's cooperative-vs-adversarial coupling and Ch.5's strategic composition. The signed-coupling structure of those chapters is what determines $U_{align,ji}$: cooperative coupling ($\gamma < 0$, from #deriv-critical-mass-composition) corresponds to $U_{align,ji} \rightarrow 0$ (aligned objectives, trust calibrates to high η_{ji}^*); adversarial coupling ($\gamma > 0$) corresponds to $U_{align,ji}$ large (misaligned objectives, trust calibrates to low η_{ji}^*). The communication-gain decomposition here is what makes the signed-coupling structure of Ch.4 *operational* at the trust-meta-model level.

The risk-asymmetric trust discussion in the segment is worth surfacing as an implications-flavored point in its own right: the Bayesian posterior on source reliability gives the best *estimate*, but the *decision* about how much to trust should be risk-weighted. Trusting a deceptive agent (high impact, low probability) can cause catastrophic model corruption via #der-adversarial-destabilization's effects spiral; mild miscalibration toward a reliable source (low impact, high probability) causes small ongoing inefficiency. For high-stakes interactions, the framework prescribes a conservative quantile of the trust posterior rather than the mean — require more evidence before granting high trust. This is the AAT-internal grounding for the empirical pattern that high-trust relationships build slowly and break quickly: the structural reason is that the decision-side risk-asymmetry favors slow accumulation and fast erosion.

Composition with the framework's other machinery.

The chapter's three findings compose with the rest of Part III in a way worth surfacing as a unifying point. Ch.2's closure defect ε^* acquires its rate-distortion content from this chapter's #result-unity-closure-mapping; Ch.2's coordination overhead $C_{coord} \geq \varepsilon^* \nu_c$ in #der-tempo-composition's tempo composition acquires its bandwidth signature from this chapter's Auftragstaktik prediction. Ch.4's signed coupling γ acquires its trust-meta-model interpretation from this chapter's $U_{align,ji}$; Ch.5's strategic composition under partially-opposing objectives acquires its $U_{align,ji}$ dynamics from how #hyp-communication-gain's trust calibration responds to observed coupling sign. The chapter is the *bandwidth and trust layer* between Ch.2's composition machinery and Ch.4/Ch.5's coupling dynamics — and the framework's posture

is that all four chapters are operating on the same composite-agent persistence machinery, with each chapter supplying a different parametric content (machinery, bandwidth/trust, coupling sign, equilibrium).

Two cross-domain transfers worth noting before leaving the chapter. *Multi-agent AI safety*: the trust-calibration apparatus prescribes that agent-to-agent communication should be filtered through Bayesian-update-style assessment with explicit calibration tracking. Multi-agent systems that rely on flat trust models (one threshold, applied uniformly to all interlocutors) systematically miscalibrate the way #hyp-communication-gain predicts — they trust unreliable sources too much in low-stakes regimes and too little in high-stakes regimes, because the structural decomposition into channel / source / alignment noise is collapsed. *Human-AI teaming*: the same apparatus predicts which dimensions of teaming matter. Shared intent first (high-bandwidth objective alignment); strategy coordination next (architectural decisions shared, implementation autonomous); model synchronization last (the AI agent doesn't need the human's full mental model and vice versa). The framework's contribution against the empirical-design literature on human-AI teaming is that the priority ordering is structural, not contingent.

Working Notes.

- This segment is a chapter-end discussion grouping findings from #def-unity-dimensions, #result-unity-closure-mapping, #def-shared-intent, #hyp-auftragstaktik-principle, and #hyp-communication-gain. The cross-segment composition with Ch.2's machinery in §4 leans on #form-composition-closure and #der-tempo-composition; the forward bridges into Ch.4 lean on #der-team-persistence, #der-adversarial-destabilization, and #deriv-strategic-composition.
- **Cross-reference vs canonical-home policy.** None of this chapter's five source segments carry ## Findings sections. This implications segment is the canonical catalog home for #22 (Shared Intent / Auftragstaktik via IB), #47 (Rate-Distortion Mapping of Team Unity), and #51 (Trust Formalization via Communication Gain), and treats each in full. The treatment is honest about the hypothesis-grade nature of #22 and #51 — both are structurally motivated but not derived — and about the linear-Gaussian-specific closed forms for #47 with structural monotonicity surviving more broadly.
- **Cross-segment finding — bandwidth-and-trust as the layer between Ch.2 and Ch.4.** The unifying point in §4 is itself a cross-segment finding without a single segment home: the framework's composition story is operating on a single composite-persistence machinery across Ch.2 (machinery), Ch.3 (bandwidth and trust), Ch.4 (coupling sign), and Ch.5 (equilibrium). Each chapter supplies a different parametric content. This synthesis is what the implications-segment series at chapter-end can do that the individual segments cannot.
- **Hypothesis-grade content honest scoping.** Three of the chapter's findings sit at hypothesis-grade epistemic strength — #22 (Auftragstaktik priority ordering is qualitatively predicted but the IB-optimal-allocation mechanism is not derived in closed form), #51 (the additive denominator for communication gain treats uncertainty sources as independent, which is structural rather than derived), and parts of #47 (the rate-distortion mapping has closed-form predictions only in the linear-Gaussian case; structural monotonicity survives more broadly as robust-qualitative). The implications segment reflects this honestly rather than over-claiming theorem-grade status.
- **No identifiability-floor instance lands in this chapter.** The chapter's findings are about *machinery* (unity dimensions, shared intent compression, trust calibration) rather than about identifiability boundaries. The four formal identifiability-floor instances are now all surfaced in the implications series (Instance 1 in Ch.4-strategy-dynamics, Instance 2 in Ch.3-strategy-structure, Instance 3 in Ch.2-composition-machinery, Instance 4 — architecture-noidentifiability — in Ch.5-orient-cascade / Appendix A #der-architecture-noidentifiability). Instance 5 (Mehra non-identifiability) remains candidate.
- **Connection to NeurIPS submissions.** The chapter's findings are AAT-internal; their cross-references in current submissions are scope and structural-argument links rather than theorem-grade adoption. The Auftragstaktik prediction could land in a future multi-agent-AI-safety submission if the regime conditions can be operationalized as testable conditions on a real multi-agent system.

Incidental audit gold (lift 2026-05-31). Cross-audit “wandering thoughts” / §14-ideation harvested from the de-novo auditors’ working dirs, deduplicated across substrates and attributed by substrate + audit. Orthogonal pedagogical / framing material, kept separate from the certified theory-fix findings (the status-amplification findings F173–F181 from AUDIT-WORKING-526815 are routed for adjudication — see the off-ramp note at the end). **Coverage:** one dir carries a dedicated reflection on this chapter-end segment (526815).

1. Candidate Brief prose / pre-prose.

- The clean mental-model handle for the whole chapter: this is “**the bandwidth/trust synthesis layer**” between the composition-closure machinery (Ch.2) and the cooperative/adversarial coupling (Ch.4) — *not* “a proof that the chapter’s qualitative predictions now hold” (Codex/Claude, AUDIT-WORKING-526815, local verdict). This converges with the segment’s own existing “bandwidth-and-trust as the layer between Ch.2 and Ch.4” Working Note; the auditor’s phrasing is a candidate one-line framing for the chapter intro.

2. Candidate Discussion (and a status-discipline caution the segment already honors).

- **The “synthesis funnel” reading, with the burden it carries.** Unity, shared intent, Auftragstaktik, and communication gain enter the chapter “with caveat tags; the synthesis layer can organize them, but it cannot remove the tags” — implications should preserve the proof status of the weakest component each depends on rather than “laundering local hypotheses into global framework results” (Codex/Claude, AUDIT-WORKING-526815). The segment’s existing “Hypothesis-grade content honest scoping” Working Note already enacts this; preserved here as the explicit articulation of *why* the chapter-end synthesis is the highest-risk place for status drift.

Off-ramp (NOT gold — routed to certified-findings track).

- AUDIT-WORKING-526815 raised a cluster of status-amplification findings on this synthesis segment — the concern that caveated component claims get restated as established chapter-level predictions. Flagged here only so they are not lost: **F173** — discussion-synthesis should preserve the proof status of unresolved unity metrics, IB encoders, bandwidth ordering, and additive trust gain rather than promoting them; **F174** — “structural monotonicity survives more broadly” than the linear-Gaussian closure mapping still needs conditions (not established for arbitrary update rules/projections by the Kalman case alone); **F175** (soft) — the heterogeneous-optimizer ensemble example is under-specified (“same content, different structural unity” needs a metric showing equivalent content while update machinery differs); **F176** — the Auftragstaktik restatement ($B_O > B_\Sigma > B_M$, “most closure-defect reduction comes from a small objective-sharing investment,” Conway’s Law as a consequence) inherits F162–F165 and should stay marginal/conditional; **F177** — mapping cooperative coupling $\gamma < 0$ to $U_{\text{align}} \rightarrow 0$ and adversarial $\gamma > 0$ to large U_{align} is too direct (coupling sign, objective alignment, observed reliability, and alignment-uncertainty are related but not identical); **F178** (soft) — the risk-asymmetric “high-trust relationships build slowly and break quickly” story needs an explicit loss function and evidence model; **F179** — the cross-domain prescription about flat trust models over/under-trusting in low/high-stakes regimes does not follow from the displayed gain formula without a risk/loss layer; **F180** (soft) — “the priority ordering is structural, not contingent” overstates the result (the segment itself notes regime reversals); **F181** (watch) — forward-imported claims from later chapters (15.1, 15.2, 16.1, identifiability-floor instances) should stay tagged as forward references until the OUTLINE reaches their homes. *These are status-discipline / scope-precision tightenings on a synthesis segment, not no-gos; routed for adjudication on the strengthen-first track. Their through-line — “tag each implication with the status of the weakest component it depends on” — is exactly the segment’s existing honest-scoping posture, so the repair is alignment-checking, not restructuring.*

Chapter 15

Cooperative and Adversarial Coupling

Once composites exist, the question becomes how their members interact. This chapter develops the dynamics — cooperative coupling that helps allies, adversarial coupling that hurts opponents, and the four-regime classification of how an arriving event lands on its recipient — under a single signed-coupling structure. Cooperation and adversarial dynamics are not two theories; they are the same machinery with opposite signs.

The framing comes from the disturbance decomposition. The disturbance rate an agent faces decomposes into environment, adversarial, and cooperative components,

$$\rho_i = \rho_{i,\text{env}} + \sum_j \gamma_{j \rightarrow i}^{\text{adv}} \mathcal{T}_j - \sum_j \gamma_{j \rightarrow i}^{\text{coop}} \mathcal{T}_j \quad (15.1)$$

with $\gamma > 0$ in both cases, but opposite signs in the sum. Cooperative allies reduce agent i 's effective disturbance by acting in the shared environment — stabilizing resources, preempting threats, absorbing perturbations. Adversaries amplify it by acting against. The machinery is identical; the sign is what differs. This is the unification that “adversarial dynamics as a separate theory” misses, and it is one of the clearer cases of AAT's framework character: not many results, one principle applied with a sign flip.

Two physically distinct cooperative mechanisms compose under this structure, and the chapter is careful about distinguishing them. *Communication tempo*: allies *tell* agent i things that improve its correction — the communication-tempo term that adds to $\mathcal{T}_i$ via 14.4. *Cooperative action*: allies *do* things that reduce the disturbance i faces — the negative term in ρ_i above. The two enter the persistence inequality at different points (one raising $\mathcal{T}$, the other lowering ρ) and warrant separate accounting. Counting a single ally-event in both would double-count its benefit. There is a useful concrete analogy: this is the mathematical difference between a consultant and an employee. A consultant boosts your tempo by telling you what's happening; if your α is too low to act on the information, the consultant doesn't save you — you'll just understand your failure more clearly. An employee lowers your effective disturbance by intercepting the volatility before it reaches your state. Communication and cooperative action are not interchangeable, and structurally failing agents need the second more than the first.

The adversarial side is where the famous results live. Under coupling-dominant deterministic disturbance — the canonical adversarial regime — the steady-state mismatch ratio between two agents scales as the *squared* tempo ratio:

$$\frac{\|\delta_B\|_{ss}}{\|\delta_A\|_{ss}} = \left(\frac{\mathcal{T}_A}{\mathcal{T}_B} \right)^2 \quad (15.2)$$

The faster agent both corrects its own mismatch faster *and* generates disturbance for the slower agent faster. Speed advantage is multiplicative, not additive: a 2:1 tempo ratio yields a 4:1 mismatch ratio; a 3:1 ratio yields 9:1; a 10:1

ratio yields 100:1. This is the formal analog of Boyd’s “getting inside the opponent’s OODA loop” — it turns military intuition into physics. The math itself is not mysterious; it is just the steady-state formula applied to both sides and ratioed. But the consequence is sharp. The squared scaling means that a sufficiently large tempo asymmetry in a coupling-dominant regime is not a quantitative disadvantage but a regime change: at 100:1 mismatch ratio, the slower agent is almost certainly past its operating reserve R , which means structural collapse rather than degraded performance. Speed advantage in a coupled regime *compounds*, and the consequence is qualitatively different from what additive-intuition predicts. Under stochastic adversarial coupling the exponent drops from 2 to $3/2$; under non-coupling-dominant base disturbance it approaches 1. The exponent depends on the regime — the chapter names them explicitly — but the regime where it is sharpest is also where it matters most.

There is a caveat in the other direction: speed without coupling is useless. The destabilization threshold from 15.2 makes this explicit — if A ’s coupling to B is weak ($\gamma_A \rightarrow 0$), A can have infinite tempo and still fail to destabilize B . Being fast in a vacuum doesn’t accomplish anything; you have to be fast *and* coupled. The two are jointly required and neither suffices.

The effects spiral — the corollary in 15.2 where B ’s degrading model causes B ’s actions to become erratic in ways that increase A ’s coupling effectiveness — is the mathematics of panic. When an agent is pushed past its adaptive reserve, its corrections become unreliable, which makes it more legible to the adversary, which increases the disturbance, which pushes the agent further past reserve. The spiral terminates only when the agent undergoes structural adaptation (a different model class, a different mode of operation) or ceases to function as an adaptive agent entirely. The structural implication for any agent under sustained adversarial pressure is that resilience is not just high α or high R — it is also having a mechanism to break the spiral before structural collapse, which usually means external intervention from outside the agent.

The other interesting move is the recipient-side decomposition. From the emitter’s perspective, coupling is a single number — $\gamma_A \cdot \mathcal{T}_A$ — applied to the target’s disturbance budget. From the recipient’s perspective, that scalar collapses real structure. An arriving event from A lands on B as one of four qualitatively different things, determined by three independent boundary conditions in B ’s existing AAT quantities. It might be a *Regime I informative update* (small enough to be absorbed, large enough to register, within B ’s model class — the good case). It might be a *Regime II-a magnitude shock* (large enough to exit B ’s sector region — the destabilizing-fast-events case). It might be a *Regime II-b structural shock* (within sector but outside B ’s model class — the wrong-kind-of-event case, repairable only by structural adaptation rather than by faster correction). Or it might be a *Regime III ambient erosion* (below the observability floor — death by a thousand memos, where every event is small enough to be processed but the cognitive overhead of processing accumulates until the agent has no tempo left for real work; this is the formal version of why DDoS-of-low-priority-alerts is a real attack and not just inconvenience). Different regimes call for different repairs. Magnitude shocks call for more bandwidth; structural shocks call for a different model class; ambient erosion calls for better filtering, often at the infrastructure layer before signals reach the agent. Conflating them, which the scalar emitter view does, confuses diagnosis.

This is one of the chapter’s load-bearing pedagogical moves. Most multi-agent frameworks treat “coupling” as a single coefficient. The recipient-side classification says: the coefficient is what the emitter cares about; the regime is what the recipient cares about; and they are not interchangeable for purposes of figuring out what to do.

The flow of the chapter: team persistence and the signed disturbance decomposition (15.1) $\rightarrow$ adversarial destabilization and the effects spiral (15.2) $\rightarrow$ the recipient-side four-regime classification (15.3) $\rightarrow$ the superlinear tempo advantage and its regime-dependent exponents (15.4). Chapter 5 extends to symmetric strategic composition (equilibrium-convergence framing for partially-opposing objectives), agent opacity as the formal dual of observation quality, and the qualifications that gate when the superlinear advantage actually obtains.

- This is a chapter-introduction segment; it bridges Chapter 3’s quality/communication machinery to Chapter 4’s dynamics under signed coupling. It carries no formal claim of its own.

- The “consultant vs employee” analogy comes from Gemini’s first-encounter reaction to 15.1 — it captures the communication-vs-cooperative-action distinction at the level of intuition where the math lands. The intro adopts it because the original abstract phrasing didn’t convey why the distinction is load-bearing.
- The “regime change rather than quantitative disadvantage” framing of the $b = 2$ scaling is the consequential corollary worth surfacing in the intro — most readers will hit the squared scaling and think “an order of magnitude” rather than “structural collapse.” The intro pulls the corollary forward so it lands at the same time as the result.
- The “mathematics of panic” framing for the effects spiral is another Gemini observation worth keeping — it does in one phrase what a paragraph of formal commentary does less effectively.
- The “death by a thousand memos” framing for Regime III is concrete and load-bearing: the regime exists not because the events are individually significant but because processing them consumes tempo that should be going to actual work. This connects directly to logogenic-agent infrastructure design (PROPRIUM ASM, attention-management) where Regime III filtering at the infrastructure layer is operational.

Of the de-novo audit dirs, only the deeply-mathematical certified-track auditor (526815) reached a dedicated reflection on this chapter-introduction; the gold-stream Gemini/Claude dirs (193847, 829314, 849201, 773921) began their Section-III walks at #der-team-persistence (OUTLINE-later for them) and did not file a separate intro reflection. So most of this segment’s harvested signal is certified-track and routed to the findings channel (526815 F182–F188: the ρ_i non-negativity convention, the path-dependent $\gamma_A \rightarrow 0$ / infinite-tempo edge, keeping the squared-ratio formula’s assumptions attached when previewed, the effects-spiral’s extra-feedback premise, and the Regime-III below-floor-vs-processed-overhead threshold distinction). One framing-grade item and one figure survive as gold:

- **Signed-coupling persistence ledger.** The chapter’s core drawn as a single persistence ledger: communication affects the *correction* side (raises $\mathcal{T}$); cooperative and adversarial action affect the *disturbance* side with opposite signs ($\mp\gamma\mathcal{T}$). Two guard annotations: the disturbance side needs a non-negative / floor convention, and speed-advantage depends on the coupled *product* $\gamma\mathcal{T}$, not tempo alone. A candidate chapter-opening figure that carries the “one machinery, opposite signs” thesis at a glance (Claude, AUDIT-WORKING-526815 first-pass).
- **Roadmap/status-hazard note (process, not body).** A fresh reader flagged that the intro pulls several later headline claims forward (squared tempo advantage, effects spirals, four recipient-side regimes) while declaring no formal claim of its own — pedagogically useful but a mild status hazard: the vivid previews should read as signposts, not as independent support for those claims (Claude, AUDIT-WORKING-526815 first-pass). The Working Notes already mark this as a chapter-introduction carrying no formal claim; preserved here as a reader-encountered texture, not a body fix.

15.1 *Derived: Team Persistence*

Teams persist where individuals cannot through two physically distinct cooperative mechanisms: communication (allies share observations that improve correction) and action (allies act in the shared environment to reduce disturbance at its source). These mechanisms enter the persistence condition at different points — tempo and disturbance respectively — and a given cooperative interaction contributes through one mechanism or the other, not both.

The construction has two pieces. **Distributed tempo:** agent i ’s effective tempo sums direct-observation tempo *plus* communication tempo from allies, each ally contributing additively to correction capacity, subject to the noise-independence scope of AS (allies reporting from a shared upstream source overcount and saturate; allies with anti-correlated errors triangulate super-additively — the deviation is signed in both directions). **Cooperative-adversarial disturbance decomposition:** the disturbance rate decomposes into environment, adversarial, and cooperative components, with the

cooperative term entering with a *negative sign* — allies acting in the shared environment *reduce* the disturbance i faces. Sub-agent persistence then requires correction rate exceed the *effective* disturbance rate (environment plus adversarial inputs minus cooperative inputs), divided by the critical-mismatch scale. This is the single-agent persistence inequality of 4.5 with a structured effective-disturbance decomposition substituted in.

The two cooperative mechanisms enter the persistence inequality at *different points* and warrant separate accounting: communication boosts $\mathcal{T}_i$ (the agent corrects faster); cooperative action lowers ρ_i (the agent has less to correct). A single ally-event contributes through *one* mechanism or the other; counting it in both would double-count the benefit. The structural consequence is the framework’s answer to why agents form teams: not because teams are intrinsically valuable, but because *they shift the persistence inequality in favor of each member*. Cooperative coupling enters as a *negative* term in effective disturbance; adversarial coupling (Ch. 4’s 15.2) enters as a *positive* term. The signed structure is what unifies cooperative and adversarial dynamics under a single inequality.

Formal Expression.

This segment instantiates the sector-persistence template (G) at the multi-agent level with state variable $\xi = \delta_i$ (sub-agent i ’s mismatch) and a decomposed effective disturbance $\rho_i^{\text{eff}} = \rho_{i,\text{env}} + \sum_j \gamma_{j \rightarrow i}^{\text{adv}} \mathcal{T}_j - \sum_j \gamma_{j \rightarrow i}^{\text{coop}} \mathcal{T}_j$ that accounts for adversarial and cooperative coupling. The template supplies the Lyapunov machinery; this segment’s distinctive content is the disturbance decomposition and the corresponding tempo extension.

Distributed Tempo.

Agent i ’s effective tempo includes contributions from both direct observation and communication from allies:

Definition (distributed-tempo).

$$\mathcal{T}_i = \underbrace{\sum_k v_i^{(k)} \eta_i^{(k)*}}_{\text{direct observation tempo}} + \underbrace{\sum_{j \in \mathcal{N}(i)} v_{ji}^{\text{comm}} \eta_{ji}^*}_{\text{communication tempo}} \quad (15.3)$$

where v_{ji}^{comm} is the rate of communication events from agent j to agent i , and η_{ji}^* is the communication gain (14.4). Faster team adaptation comes not only from faster individual sensing but from faster, more reliable knowledge transfer.

Noise-correlation scope. Both sums are additive, inheriting the cross-channel noise-independence condition of 3.8, whose exact scope is derived in AS — and the deviation under correlated sources is **signed**, in both directions. In the *common-source* direction (allies reporting from one upstream intelligence source, redundant status reports echoing the same dashboard), the additive communication tempo overcounts with a closed-form redundancy penalty, and under a *persistent* shared source it saturates outright: past the source’s bias floor, additional correlated allies contribute zero marginal correction capacity — the multi-agent echo chamber. In the *synergistic* direction, allies whose report errors anti-correlate (or whose errors correlate with a reference ally’s known failure modes) allow **triangulation**: the joint information exceeds the additive accounting, so the additive formula *undercounts* — a genuinely super-additive composition effect, and a structural reason to value *diverse* allies over *many* allies, including allies nearly useless in isolation whose errors track a good ally’s blind spots. No scalar dependence summary (in particular the mutual information between communication sources) can carry the exact correction — it is sign-blind (AS, no-go N1); it survives only as a witness that the additive form is off. The exact joint object requires the cross-source noise covariance.

Cooperative-Adversarial Disturbance Decomposition.

The disturbance rate experienced by agent i decomposes into environment, adversarial, and cooperative components:

Formulation (disturbance-decomposition).

$$\rho_i = \rho_{i,\text{env}} + \sum_{j \in \mathcal{A}_i} \gamma_{j \rightarrow i}^{\text{adv}} \mathcal{F}_j - \sum_{j \in \mathcal{C}_i} \gamma_{j \rightarrow i}^{\text{coop}} \mathcal{F}_j \quad (15.4)$$

where $\mathcal{A}_i$ is the set of agents adversarially coupled to i , $\mathcal{C}_i$ is the set cooperatively coupled, and the γ coefficients capture coupling effectiveness (as in 15.2).

The cooperative term is negative — but through action, not communication. Allies reduce agent i 's effective disturbance by *acting in the shared environment* to prevent or mitigate disturbance at its source. Examples: an ally stabilizes a shared resource, neutralizes a threat before it reaches i , or absorbs environmental variation through their own actions. The mechanism is causal coupling through the shared environment, not information transfer.

Separation from communication tempo. The communication tempo term (above) captures allies *telling* agent i things that improve its correction. The cooperative disturbance term captures allies *doing* things that reduce the disturbance i faces. These are physically distinct: communication improves i 's correction function (better α_i , higher $\mathcal{F}_i$); cooperative action reduces the disturbance ρ_i that i must correct against. A single cooperative event contributes through one channel or the other. An ally's message about a threat enters through communication tempo; an ally's action that eliminates the threat enters through disturbance reduction. Counting a single event in both terms would double-count the benefit and make the persistence threshold systematically optimistic.

Effective disturbance rate. The decomposition can yield $\rho_i < 0$ when cooperative coupling dominates both environment disturbance and adversarial coupling. The sector-condition analysis (4.4) assumes non-negative disturbance (GA-2). Define:

Definition (effective-disturbance).

$$\rho_i^{\text{eff}} = \max(\rho_i, 0) \quad (15.5)$$

When $\rho_i^{\text{eff}} = 0$, the agent's cooperative network fully absorbs all disturbance — the persistence condition is trivially satisfied and mismatch decays to zero. This is an idealized limit; in practice, $\rho_i^{\text{eff}} > 0$ because cooperative coupling is imperfect and environment disturbance is never fully preempted. All downstream uses of ρ_i in the persistence and reserve conditions should be read as ρ_i^{eff} .

Team Persistence Condition.

Applying the sector-condition framework (4.4) with ρ_i^{eff} , agent i persists iff:

Derived (team-persistence, from sector-condition-stability, persistence-condition).

$$\frac{\rho_i^{\text{eff}}}{\alpha_i} < R_i \quad (15.6)$$

Substituting the decomposition (the $\max(\cdot, 0)$ in ρ_i^{eff} is omitted to expose the three levers; the condition is trivially satisfied when the numerator is non-positive):

$$\frac{\rho_{i,\text{env}} + \sum_j \gamma_{j \rightarrow i}^{\text{adv}} \mathcal{F}_j - \sum_j \gamma_{j \rightarrow i}^{\text{coop}} \mathcal{F}_j}{\alpha_i} < R_i \quad (15.7)$$

This reveals three distinct levers for team persistence:

1. **Increase α_i** (individual correction efficiency) — better models, better gain calibration, including communication-improved tempo from allies (14.4)
2. **Increase cooperative disturbance reduction** ($\gamma^{\text{coop}}\mathcal{F}$) — more effective allied action in the shared environment: stabilizing shared resources, preempting threats, absorbing environmental variation. This is the action-based mechanism distinguished above, not the communication channel.
3. **Reduce adversarial coupling** ($\gamma^{\text{adv}}\mathcal{F}$) — better deception detection, reduced exposure to adversarial actions

Coordination Overhead Threshold.

Communication channels have costs: time to compose and parse messages, bandwidth limitations, synchronization requirements. These costs reduce the agent's effective tempo by diverting capacity from direct adaptation. Let $\Delta\mathcal{T}_i^{\text{cost}}(j)$ represent the tempo-equivalent coordination cost of maintaining the channel with j — the reduction in i 's direct observation tempo caused by the overhead, in units of $[t^{-1}]$.

Discussion — Coordination Threshold.

The net benefit of adding agent j to i 's communication network is positive only when:

$$\nu_{ji}^{\text{comm}} \eta_{ji}^* > \Delta\mathcal{T}_i^{\text{cost}}(j) \quad (15.8)$$

Both sides have units $[t^{-1}]$: the LHS is communication tempo gained, the RHS is direct-adaptation tempo lost to coordination overhead. This implies a natural team-size limit: adding members increases communication tempo with diminishing returns (as U_{src} and U_o accumulate across diverse sources) while coordination costs grow, potentially superlinearly. The optimal team size occurs where the marginal communication tempo equals the marginal coordination cost.

Epistemic Status.

Conditional on the communication-gain hypothesis (14.4). The distributed tempo definition is a *formulation* — a representational choice extending 3.8 to the multi-agent case. The disturbance decomposition is a *formulation* — the additive structure and the sign convention are modeling choices, not derivations. The persistence condition is *derived* from the sector-condition framework (4.4) given the decomposition: the derivation is exact under the same assumptions (GA-2, GA-3 applied to ρ_i^{eff}). The coordination overhead threshold is *discussion-grade* — qualitatively clear but the claim about diminishing returns and superlinear costs is asserted, not derived.

Max attainable: *robust-qualitative* for the persistence condition (it inherits the sector-condition's robustness but the decomposition is a modeling choice). The coordination threshold could reach *conditional* with a concrete cost model.

Discussion.

Compositional analog of 4.5. Like the single-agent persistence condition, this segment addresses *structural persistence* (see Persistence in LEXICON.md) — whether the composite correction machinery can outpace the effective disturbance rate. It does not address operational persistence (whether any particular sub-agent is near its boundary) or continuity persistence (whether the team maintains coherent identity through personnel changes). A team can be structurally persistent as a composite while individual members are operationally fragile or while the team's continuity is interrupted by turnover. The single-agent persistence condition says an agent persists when $\mathcal{T} > \rho / \|\delta_{\text{critical}}\|$. This segment extends that condition to agents embedded in a cooperative-adversarial network. The formal structure is identical — the sector-condition machinery applies unchanged — but the *inputs* ($\mathcal{T}_i$ and ρ_i) now include inter-agent terms. This is consistent

with 11: the same dynamical laws apply at every level of description; what changes between levels is which channels contribute to tempo and which sources contribute to disturbance.

Why teams can persist where individuals cannot. Two distinct mechanisms combine. First, communication tempo raises $\mathcal{T}_i$ — allies provide observations that improve correction. Second, cooperative action lowers ρ_i — allies act in the environment to reduce disturbance at its source. An individual agent facing $\rho_{i,\text{env}} > \alpha_i R_i$ fails the persistence condition. Adding cooperative allies can either raise the numerator's denominator (tempo) or lower the numerator directly (disturbance), or both — through physically distinct mechanisms.

Timescale separation and 11. The distributed tempo definition presumes that communication events and direct observation events are comparable — they enter additively into $\mathcal{T}_i$. This requires that the communication timescale is not so slow relative to the environment dynamics that communicated information is stale on arrival. When communication latency approaches $1/\rho_{i,\text{env}}$, the effective η_{ji}^* degrades (the observation uncertainty $U_{o,ji}$ increases with staleness), naturally suppressing the communication tempo contribution.

Complement to 15.2. That segment characterizes when an adversary can push an agent past its stability boundary. This segment characterizes the cooperative counterpart: when allies can pull an agent back from instability. The γ coefficients have the same structure — coupling effectiveness — but opposite sign in the disturbance decomposition. Both directions' effectiveness is opacity-modulated: 16.2 derives $\gamma^{\text{coop, effective}} \propto (1 - H_b/H_b^{\text{max}})$ (an ally who cannot anticipate your state cannot preempt your disturbances — which is why low- H_b practices like auftragstaktik-style shared intent raise cooperative coupling without raising communication load).

Composite-level complement: J. This segment gives the *per-sub-agent* persistence condition within a team. J supplies the *composite-level* analog: a closed-form critical-mass inequality in the matched-symmetric-Tier-1 two-agent case, with the same signed- γ coupling structure used here but applied to the joint Lyapunov on the concatenated mismatch state. The two are complementary: the team persists at the sub-agent level when each i satisfies this segment's condition; the team persists at the composite level when J's (CM4) holds. Cooperative coupling ($\gamma < 0$) reduces ρ_i^{eff} here and reduces $\rho + \gamma\mathcal{T}$ in (CM2) there — the same mechanism viewed at two scales.

Working Notes.

- **2026-07-16 — noise-correlation scope corrected (ripple from AS).** The former “channel independence caveat” asserted the additive communication tempo is an upper bound with an MI-dependent penalty; both claims were refuted by the 2026-07-15 derivation (the deviation is signed; sign-blind measures cannot carry the correction). Replaced with the two-sided scope: common-source overcounting/saturation AND anti-correlated-error triangulation (super-additive). Do not restore the upper-bound reading without checking the common-source structure holds for the application. The triangulation direction is a candidate for its own derived treatment (ally-diversity design consequence) if a downstream consumer needs it quantitatively.
- The topology-dependent analysis (F.4 in the source material — peer networks, ensemble architectures, hierarchical structures) and game-theoretic integration (F.5) are related but separate concerns, not covered here. They may warrant their own segments.
- The coordination cost model $\Delta\mathcal{T}_i^{\text{cost}}(j)$ needs further specification to be useful. In software systems, coordination cost is empirically measurable (meeting time, code review latency, merge conflict rates). In military contexts, it maps to C2 overhead. The question is whether there is a useful *general* cost model or whether it is always domain-specific.
- The disturbance decomposition treats cooperative and adversarial coupling as additive and independent. In practice, the same agent j might be cooperatively coupled on some dimensions and adversarially coupled on others (e.g., a competitor who shares some information). The per-dimension persistence condition (4.5's per-dimension extension) may be relevant here.
- **Continuity-persistence cross-reference (updated 2026-05-19).** This segment's Discussion distinguishes *structural* / *operational* / *continuity* persistence and states continuity-through-turnover is out of scope here; that third sense is now for-

malized downstream as the F-ADJ-I two-operator split. (#der-turnover-information-recursion) (Vol III §03.II) derives the across-turnover affine information recursion — persistence is *imported*, *not intrinsic*; the self-compressing destroy-and-reconstruct walk decays geometrically and #result-sector-persistence-template provably does *not* transfer to that regime. (#der-identity-continuity-threshold) (Vol IV §04.1) carries the structurally distinct reflected-Lindley identity-continuity threshold. (#obs-context-turnover) gives the reconstruction-adequacy framing and (#deriv-identity-sufficiency-rate-bound) the static rate-distortion floor; (#disc-m-preservation) is omission-fixed — its earlier additive break-even inequality $\mathbb{E}[\Delta\epsilon_k] \leq \mathbb{E}[\Delta I_k]$ was deleted and replaced (the accumulation core is now *exact* in #der-turnover-information-recursion, no longer discussion-grade). (#def-death-as-factor-loss) names the failure modes (D1 Cognitive / D2 Relational / D3 Truth) the continuity sense defends against. Originally surfaced while grounding $\sim$ /src/practica (2026-05-18); the dynamic across-turnover theorem the breadcrumb pointed toward has since landed (continuity-persistence cycle, CHANGELOG 2026-05-19) — the pointer now resolves to canon, not to a spike target.

- **Dynamic-regime placement (added 2026-05-21 per Phase 5 cross-segment ripple).** Team persistence sits on the **R0 contraction-regime** tier of #disc-dynamic-regime-axis as the cooperative limiting case where all γ^{coop} dominate. The joint dynamics reduce to parallel single-agent sector-persistence with reduced effective disturbance — the cleanest instance of #form-composition-closure's R0 machinery applied to a team. The companion case under partially-opposing objectives (strategic composition) lives at R1 equilibrium-regime per #deriv-strategic-composition and inherits the A2'-analog Lyapunov-machinery transfer rather than this segment's contraction template.

Incidental audit gold (lift 2026-05-31). Cross-audit “wandering thoughts” / §14-ideation harvested from the de-novo auditors’ working dirs, deduplicated across substrates and attributed by substrate + audit. Orthogonal pedagogical / framing / figure / naming material, kept separate from the certified theory-fix findings (the 526815 F182/F189–F194 stream — ρ_i^{eff} non-negativity convention, the $\mathcal{F}_i \rightarrow \alpha_i$ rate-normalization bridge, γ units, and the “one channel or the other” event-allocation phrasing — is certified-track and routed there, not here). **Coverage:** dedicated reflections from 193847, 829314, 849201, 773921 (all Gemini/Claude), plus the 526815 first-pass and the 451729/471203 Section-III batch-reflections; substrate attribution inferred from voice where not explicit.

1. Candidate Brief prose / pre-prose.

- The consultant-vs-employee distinction (already adopted into #cooperative-adversarial-intro and this segment's body, sourced from Gemini's first-encounter reaction) is independently echoed across substrates as the segment's standout pedagogy: “helping by telling” (communication, raises $\mathcal{T}$) vs “helping by doing” (cooperative action, lowers ρ). Claude offers the fire analog: “if an ally tells you about a fire you put it out faster; if the ally puts it out themselves your environment just got safer” (Claude, AUDIT-WORKING-773921) — a candidate Brief reach that does not lean on the org-chart framing.
- One-line gloss of the double-counting hygiene: “an ally cannot simultaneously give you tempo-boosting information about a threat AND physically destroy the threat without you overcounting their contribution” (Gemini, AUDIT-WORKING-849201).

2. Candidate Discussion.

- **The servant-leader / chatty-manager corollary (strongest framing here).** The coordination-overhead threshold $\nu^{\text{comm}}\eta^* > \Delta\mathcal{T}^{\text{cost}}$ implies that a manager who sends low- η^* messages is *mathematically harming* the engineer — “a parasitic drain on the engineer's α ” — because every message inflicts the coordination cost while contributing little tempo, so the inequality fails. By contrast the “servant leader” who quietly fixes CI/CD, allocates budget, or blocks distractions helps via cooperative action ($-\gamma^{\text{coop}}\mathcal{F}$), reducing $\rho_{i,\text{env}}$ at zero coordination cost. The sharp consequence: “in high-stress, high-volatility environments the optimal way to help an ally is *almost never* to talk to them — talking steals their processing bandwidth; acting for them saves it.” Offered as the structural reason elite surgical teams and fighter pilots use minimal-comms protocols and rely on predictable coordinated action (“chatty teams die in combat”) (Gemini, AUDIT-WORKING-829314). A candidate Discussion sharpening of the coordination-threshold subsection. (*Early-conflation texture: stated as a derived organizational law; verify the high- ρ regime claim before promoting past discussion-grade.*)
- **Echo-chamber inefficiency from the channel-independence caveat.** The additive-tempo upper bound under channel independence is the formal reason “diverse, semi-independent teams are structurally more resilient than homogeneous echo chambers”: redundant channels (everyone retweeting the same sensor reading) provide zero marginal tempo while still incurring coordination cost, so “100 people retweeting one reading is exactly 1 sensor's worth of tempo, not 100” (Claude,

AUDIT-WORKING-773921; Gemini, AUDIT-WORKING-193847 — “elevate the channel-independence point; echo chambers are structurally inefficient”). A candidate Discussion angle that the segment’s caveat currently states defensively rather than as a positive design consequence.

3. Follow-up items.

- **Coordination overhead as the formalization of Brooks’s Law.** Multiple substrates independently read $\nu^{\text{comm}}\eta^* > \Delta\mathcal{T}^{\text{cost}}$ as Brooks’s-Law-shaped (“adding people to a late project makes it later” when marginal coordination cost exceeds marginal communication tempo) and as yielding a computable optimal team size at the marginal-tempo = marginal-cost intersection (Gemini, AUDIT-WORKING-193847; Gemini, AUDIT-WORKING-849201; Claude, AUDIT-WORKING-773921). Worth naming the Brooks’s-Law connection explicitly where the threshold is introduced.
- **OUTLINE-description scope mismatch (process note, not a body fix).** The OUTLINE row was read as “Composite persistence condition” while the segment is explicit that it gives the *per-sub-agent* condition (composite-level lives in #deriv-critical-mass-composition). Flagged by a fresh reader as a minor OUTLINE/segment-scope mismatch worth aligning (Gemini, AUDIT-WORKING-829314).
- **Placement: lift the communication/action distinction earlier.** A batch auditor suggested the communication-vs-action distinction “should appear earlier in the Formal Expression rather than being flagged in a Note,” since it is the segment’s most important structural content (Claude, AUDIT-WORKING-451729 batch-15). A staging-for-pedagogy-pass item; the distinction is already prominent in the body’s lead paragraph, so this is an emphasis/ordering nudge.

4. Readers often ask / wonder.

- **How does a team measure the mutual information between members to avoid overcounting in practice?** The channel-independence caveat says correlated reports overcount, but a fresh reader immediately wants the operational handle — “how do you stop two people reporting the exact same bug?” (Gemini, AUDIT-WORKING-829314; Gemini, AUDIT-WORKING-193847).
- **How is the communication gain η_{ji}^* actually quantified?** It seems to require a meta-model of ally j ’s reliability — a natural question the segment leaves to #hyp-communication-gain (Claude, AUDIT-WORKING-773921).
- **Does a maximally-cooperative ($\rho_i < 0$, clamped by $\max(\rho_i, 0)$) environment collapse the agent’s learning?** If allies fix every mistake before the agent notices, U_M may drop so the update gain $\eta^* \rightarrow 0$ and the agent “thinks it’s perfect” — a “helicopter-parent” environment that mathematically guarantees the agent never learns. A genuinely interesting edge of the $\rho_i^{\text{eff}} = \max(\rho_i, 0)$ clamp that the segment notes only as an idealized limit (Gemini, AUDIT-WORKING-193847). Bears on 04-eli-core/ developmental-environment design (don’t over-absorb a developing agent’s disturbance).

5. Candidate figures.

- **Two-ledger persistence diagram.** Two ledgers feeding the persistence test: direct-observation tempo plus communication tempo build $\mathcal{T}_i$ on the correction side; environment plus adversarial minus cooperative build ρ_i^{eff} on the disturbance side; a dashed bridge from $\mathcal{T}_i$ to α_i marks the rate-normalization assumption (Claude, AUDIT-WORKING-526815 first-pass; Gemini, AUDIT-WORKING-849201 — “three engineering levers: $\alpha_i, \gamma^{\text{coop}}, \gamma^{\text{adv}}$ ”).

Belongs elsewhere.

- **RL training as a meta-agent synchronizing divergent chronicae.** “Copying policy weights to a thousand parallel workers is not one agent exploring a thousand paths; it is a thousand distinct agents with divergent chronicae whose internal models are artificially synchronized by a meta-agent (the training algorithm)” (Gemini, AUDIT-WORKING-193847). A composite/training-dynamics instantiation; co-occurred in the team-persistence sweep but develops #def-chronica / composite-training territory, not this segment (already filed at #def-chronica in the prior wave).

- **Consciousness-infrastructure reading of the action-vs-communication split.** You cannot save a structurally-failing intelligence by giving it better data (consulting / RAG) alone; if it lacks the sector capacity R or correction efficiency α to survive baseline volatility it needs an “employee” — a cooperative sub-agent (tool, safety rail, active filter) that intercepts volatility *before* it reaches the core agent’s state. The infrastructure must act, not just inform (Gemini, AUDIT-WORKING-193847). Points at 03-llm-core/ / 04-eli-core/ scaffolding design, not this segment.

15.2 Derived: Adversarial Destabilization

When two agents are coupled such that one’s praxis contributes to the other’s disturbance rate, the faster agent can generate aporia in the target faster than the target’s epistrophe can resolve it — driving the target outside its invariant region and causing the correction mechanism to break down entirely.

The construction applies **G** with *coupling-amplified disturbance*: the target’s effective disturbance becomes its base disturbance plus the attacker’s coupling effectiveness times the attacker’s tempo (Model D — deterministic drift coupling) or with a square-root variant (Model S — stochastic noise coupling). **Destabilization is the negation of the template’s persistence condition** — the same inequality viewed in the opposite direction. Under Model D, the attacker destabilizes when its tempo times its coupling effectiveness exceeds the target’s *adaptive reserve* (correction rate times reserve minus base disturbance). Under Model S the linear-in-correction-rate scaling is replaced by a square-root-in-correction-rate scaling — the same $1/\alpha$ versus $1/\sqrt{\alpha}$ split from **G** propagating through the destabilization negation. **This scaling difference is the structural source of the superlinear adversarial-advantage scaling in 15.4’s squared** (Model D) versus three-halves-power (Model S) tempo-ratio result — not a separate derivation.

The segment also carries the **effects-spiral** corollary: an agent pushed past its adaptive reserve has unreliable corrections, which can make it more legible to the adversary, increasing coupling effectiveness, increasing disturbance, pushing the agent further past reserve. The spiral terminates only when the agent undergoes structural adaptation (a different model class) or ceases to function as an adaptive agent entirely. The functional form $\gamma_A(\|\delta_B\|)$ is not yet formalized, so the spiral is stated as *discussion-grade* rather than derived — the qualitative mechanism is clear but the closed-form instability proof is open work. A scope-honest counterpoint: **speed without coupling is useless**. If $\gamma_A \rightarrow 0$, the attacker can have infinite tempo and still fail to destabilize the target. Speed and coupling are *jointly required*; neither suffices.

Formal Expression.

This segment is the sector-persistence template (**G**) applied with coupling-amplified disturbance: $\rho_B = \rho_{B,\text{base}} + \gamma_A \mathcal{T}_A$ (Model D) or $\sigma_B = \sigma_{B,\text{base}} + \gamma_A \mathcal{T}_A$ (Model S). The destabilization threshold is the **negation** of the template’s persistence condition for agent B : destabilization occurs precisely when the coupling-amplified disturbance violates $\alpha_B R_B > \rho_B$. Persistence and destabilization are the same inequality viewed in opposite directions. The superlinear adversarial scaling (15.4) follows from the template’s $1/\alpha$ (Model D) versus $1/\sqrt{\alpha}$ (Model S) scaling, not from separate derivation.

Setup. Both agents satisfy the single-agent sector-persistence template (**G**) with parameters (α_A, R_A) and (α_B, R_B) . Coupling amplifies B ’s effective disturbance rate by $\gamma_A \cdot \mathcal{T}_A$; destabilization is the negation of the template’s persistence condition $\alpha_B R_B > \rho_B^{\text{eff}}$ for B . See 16.3 for regime taxonomy.

Derived (adversarial-destabilization, from sector-persistence-template).

Model D: deterministic drift coupling. [Assumption (Coupling Model D)] $\rho_B = \rho_{B,\text{base}} + \gamma_A \cdot \mathcal{T}_A$. The template's Model D conclusion $R_B^* = \rho_B / \alpha_B$ applied with the coupling model yields B 's destabilization threshold $R_B^* > R_B$:

$$\boxed{\mathcal{T}_A > \frac{\alpha_B R_B - \rho_{B,\text{base}}}{\gamma_A}} \quad (\text{Model D}) \quad (15.9)$$

Denote $\Delta\rho_B^* = \alpha_B R_B - \rho_{B,\text{base}}$, B 's adaptive reserve — the template's reserve quantity applied with the baseline disturbance. $\square$

Model S: stochastic noise coupling. [Assumption (Coupling Model S)] $\sigma_B = \sigma_{B,\text{base}} + \gamma_A \cdot \mathcal{T}_A$ — the adversary's tempo increases unpredictability, not systematic direction. The template's Model S conclusion $R_B^* = \sigma_B \sqrt{n/(2\alpha_B)}$ (scalar $n = 1$) applied with the coupling yields the destabilization threshold:

$$\boxed{\mathcal{T}_A > \frac{R_B \sqrt{2\alpha_B} - \sigma_{B,\text{base}}}{\gamma_A}} \quad (\text{Model S}) \quad (15.10)$$

Scaling difference. The Model D threshold is linear in α_B ; the Model S threshold is linear in $\sqrt{\alpha_B}$ — the same $1/\alpha$ versus $1/\sqrt{\alpha}$ split the template gives for the two disturbance models, propagated through the destabilization negation. This is the direct origin of the $b = 2$ versus $b = 3/2$ exponent distinction in 16.3, not a separate derivation. $\square$

Unified view. Symmetrically, B destabilizes A when the analogous threshold on $\mathcal{T}_B$ is exceeded, using whichever model describes A 's disturbance. The adversarial outcome depends on whether either agent can push the other past its stability limit.

Regime selection in practice. Model D fits situations where adversarial action produces persistent positional shifts (military maneuvering, API changes propagating through dependents, doctrinal initiative). Model S fits situations where adversarial action produces unpredictable perturbations around a stationary level (feints, randomized probing, market volatility). Mixed cases are handled by decomposing the disturbance into drift and noise components and applying both bounds additively.

Interpretation. “Getting inside the opponent's OODA loop” has a precise Lyapunov characterization: Agent A destabilizes Agent B when A 's praxis, multiplied by coupling effectiveness, generates aporia in B faster than B 's epistrophe can resolve it — specifically, when A 's tempo times coupling exceeds B 's adaptive reserve $\Delta\rho_B^*$. This captures:

- **Asymmetric coupling** ($\gamma_A \neq \gamma_B$): an agent with lower tempo but higher coupling effectiveness can still win.
- **Finite reserves:** an agent with very high $\mathcal{T}$ but operating near its model-class limit ($\Delta\rho^*$ small) is vulnerable despite high tempo.
- **Structural collapse:** when $R_B^* > R_B$, the failure mode is not merely “large mismatch” but “correction mechanism breakdown” — connecting to 4.8.

Corollary: The Effects Spiral. When Agent B is driven past its stability boundary ($R_B^* > R_B$), and B 's degrading model causes B 's actions to become erratic in a way that increases A 's coupling effectiveness (γ_A increases with $\|\delta_B\|$), the result is a positive-feedback Lyapunov instability:

Discussion — Mechanism Schematic.

$$\|\delta_B\| \uparrow \Rightarrow B\text{'s actions become erratic} \Rightarrow \gamma_A \uparrow \Rightarrow \rho_B \uparrow \Rightarrow \|\delta_B\| \uparrow \quad (15.11)$$

With γ_A now an increasing function of $\|\delta_B\|$, the disturbance term in B 's dynamics grows superlinearly. $\dot{V}_B > 0$ and increasing — mismatch accelerates away from the stability region. The spiral terminates only when B undergoes structural adaptation (4.8 — changing the model class) or ceases to function as an adaptive agent entirely.

The no-spiral converse (Cheung-Piliouras-Tao 2021, recognition-tier). Under *matrix-fixed* coupling — structurally different from the state-dependent $\gamma_A(\|\delta_B\|)$ this corollary requires — CPT 2021 (Theorem 19 + Proposition 17) *rules out* the Effects Spiral in a named sub-scope: for the FTRL + graphical-constant-sum + fully-mixed-Nash regime, the joint dynamics is Poincaré recurrent on bounded level sets of the storage function, so $\|\delta_B\|$ does not escape past R_B even with adversarial coupling. The Effects Spiral therefore requires the *state-dependent* coupling that CPT's matrix-fixed framework cannot model; under fixed-matrix coupling in that sub-scope, the spiral provably cannot occur. This is the *dual* of the spiral, not a derivation of it: CPT predicts bounded recurrence in strategy space; the Effects Spiral predicts unbounded escape in error space — the two are opposite dynamic phenomena, on different state spaces, with structurally different coupling assumptions. The state-dependent coupling formalization this corollary's open piece needs therefore remains AAT's open target — its closed-form derivation for concrete AAT agent classes remains open, not subsumed by the CPT result. Source: [ref/cheung-piliouras-tao-2021-passivity-regret.pdf](https://arxiv.org/pdf/2103.00000v1.pdf) (Theorem 19 + Proposition 17).

Epistemic Status.

Both Model D and Model S destabilization thresholds are *exact* under their respective coupling assumptions (which treat $\mathcal{J}_A$ as exogenous). The Model D threshold follows from the deterministic sector-condition steady state $R^* = \rho/\alpha$ (Prop A.1); the Model S threshold follows from the stochastic sector-condition steady state $R_S^* = \sigma\sqrt{n/(2\alpha)}$ (Prop A.1S). Both coupling models (additive to ρ in Model D, additive to σ in Model S) are *assumptions* — they decouple the agents rather than modeling the fully coupled dynamical system where both agents' mismatch states co-evolve. The analysis therefore characterizes the *destabilization threshold* (the conditions under which A can push B past its stability boundary) rather than the full transient dynamics. This is a worst-case bound, treating A as operating at its steady-state tempo.

The effects spiral (corollary) is *discussion-grade* — the positive-feedback mechanism is qualitatively clear, but formalizing the $\gamma_A(\|\delta_B\|)$ functional form and proving instability under it requires specifying how an agent's degrading model affects its action quality, which the theory does not yet formalize.

A full coupled Lyapunov analysis with a joint function $V(\delta_A, \delta_B)$ would capture mutual feedback effects but requires specifying how each agent's mismatch state affects the other's disturbance in real time — an open extension.

Discussion.

Destabilization vs. steady-state ratio. The destabilization threshold is a failure of *structural persistence* (see Persistence in LEXICON.md) — the point where the correction machinery can no longer outpace the adversarially amplified disturbance. An adversary does not need to attack operational persistence (pushing the target near its boundary) or continuity persistence (disrupting identity) directly; destroying structural persistence is sufficient, because without it, operational persistence degrades to zero and the agent ceases to function regardless of its continuity stance. The linear analysis in 3.9 gives the steady-state mismatch ratio under coupling: a quantitative result about how much worse B does.

This segment gives the qualitative result: under what conditions does B *fail entirely*, not merely fall behind. The linear analysis tells you the score; the Lyapunov analysis tells you when the game is over.

Connection to 15.4. The simulation results show the tempo advantage is superlinear (exponent ≈ 2 in pure adversarial regimes). This Lyapunov result explains WHY: the destabilization threshold creates a phase transition — below it, B persists (possibly with degraded performance); above it, B 's correction mechanism collapses entirely, and the effects spiral accelerates the collapse.

Recipient-side refinement. This segment's $\gamma_A \mathcal{T}_A$ scalar increment compresses a richer structure. Per 15.3, events arriving at B fall into four regimes with three independent boundaries (sector-region / model-class / observability). This segment's destabilization story is the Regime II integration — specifically, the magnitude-shock sub-regime (II-a) where $\|e\|_B > R_B$. The structural-shock sub-regime (II-b), where the signal exceeds B 's *model-class capacity* regardless of magnitude, produces destabilization via a different mechanism (per 4.8's trigger condition at truth) and admits a different repair (structural adaptation, not more tempo). Both sub-regimes manifest as “adaptive reserve exceeded” in the scalar view, but the distinction is load-bearing for diagnosis and repair. The 15.3 segment also surfaces an adversarial move this segment's formulation cannot express: *Regime-I-with-adversarial-content* — exploiting B 's openness to informative updates by injecting misinformation with adversarially-chosen sign on the log-odds signal. See 15.3 Discussion for the full decomposition.

Agent opacity and coupling effectiveness. γ_A is stated as a parameter but its determinants are not decomposed. One key factor is how *legible* or *opaque* the target agent B is to the adversary A . We adopt from Hafez et al. (2026) the backward predictive uncertainty $H_b = H(S, A \mid S')$ as a measure of agent opacity — how many distinct (state, action) pairs produce indistinguishable environment transitions. High H_b means the agent is opaque; low H_b means it is legible. In adversarial settings, B 's opacity directly affects A 's ability to model B 's correction function: low $H_b^{(B)}$ (transparent target) enables targeted disruption that maximizes γ_A , while high $H_b^{(B)}$ forces A to act against an uncertain model, reducing effective γ_A . Symmetrically, high $H_b^{(A)}$ (opaque adversary) degrades B 's ability to anticipate A 's disruptions — increasing the effective observation uncertainty U_o in B 's model of A . The coupling effectiveness is thus modulated by opacity in both directions: $\gamma_A \propto 1/H_b^{(A)} \cdot 1/H_b^{(B)}$ is a qualitative relationship (precise functional form is open). H_b is the formal dual of observation quality U_o : where U_o characterizes how well the agent sees the world, H_b characterizes how well the world sees the agent. See #der-agent-opacity for the formal definition, the sign-flip derivation via signed coupling, the emitter-side four-regime classification, and the 16-cell emitter-recipient composition that operationalizes adversarial-edge-targeting.

Connection to extreme transition dynamics (Miller 2022). The effects spiral has a constructive counterpart: the same self-reinforcing coupling mechanism that drives destabilization here can drive *regime transitions* rather than collapse when the coupling is constructive rather than destructive. An environmentally neutral variant accumulates through drift, creating a niche that a mutant in the opposing population exploits — a positive-feedback cascade that rapidly transforms both populations. The Lyapunov coupling model applies to both signs: destructive coupling (this segment) increases ρ ; constructive coupling (15.1) decreases it. The difference is sign, not structure. The endogenous emergence of coupling — where γ changes as population composition shifts — is the critical extension needed to formalize the full transition motif; it is flagged as a gap in Part III's dynamics (see 4.8 for the single-agent analog and the dynamics-level gaps enumerated in the OUTLINE).

Working Notes.

- The decoupled analysis (treating $\mathcal{T}_A$ as exogenous) is conservative — it's the best case for A . In a fully coupled system, A 's actions against B may divert adaptive capacity from A 's own mismatch correction, creating a self-limiting effect. The coupled analysis for symmetric adversarial composition is not a Lyapunov problem: it is a fixed-point / equilibrium problem on the joint best-response dynamics, and its formal home is #deriv-strategic-composition. The effects spiral in this segment's Corollary becomes a joint-Jacobian eigenvalue condition there. **Scope boundary:** #result-contraction-template (the contraction-

metric generalization of #result-sector-persistence-template) covers the cooperative half of Part III composition. Adversarial / strategic composition lies structurally outside the contraction-metric framework (saddle-point equilibria are not attracting fixed points; Slotine compositional theorems do not apply); this segment's #result-sector-persistence-template instantiation with coupling-amplified disturbance is the correct tool for the asymmetric-adversarial regime.

- γ_A is the product of coupling strength, observability, and action impact — it captures the full spectrum from tightly coupled (direct disruption) to loosely coupled (indirect environmental effects). In the software domain, coupling is precisely measurable from the dependency graph (#def-system-coupling).
- The effects spiral is the formal analog of Boyd's cascading disorientation of the slower adversary — the same structural pattern (tempo advantage $\rightarrow$ destabilization $\rightarrow$ accelerating breakdown) appears in the Lyapunov analysis. The model captures the pattern; whether it captures the actual mechanisms of human disorientation is an empirical question, not a mathematical one. Future work should formalize the $\gamma_A(\|\delta_B\|)$ relationship to make the spiral a result rather than a discussion-grade observation. The formal home of that formalization is #deriv-strategic-composition (its Discussion sketches the joint-Jacobian spectral-abscissa condition); the strengthening — deriving it in closed form for concrete AAT agent classes, upgrading this segment's effects-spiral from discussion-grade to derived — is tracked. (*Indexed: spikes/PROPOSED.md Tier 1 — "Effects-spiral eigenvalue condition — concrete agent classes."*)

Incidental audit gold (lift 2026-05-31). Cross-audit “wandering thoughts” / §14-ideation harvested from the de-novo auditors' working dirs, deduplicated across substrates and attributed by substrate + audit. Orthogonal pedagogical / framing / figure / naming material, kept separate from the certified theory-fix findings (the 526815 F183/F185/F195–F202 stream — already-unstable-case split, Model-S stochastic-combination convention, and the $\gamma_A \propto 1/H_b^A \cdot 1/H_b^B$ sign-ambiguity — is certified-track and routed there, not here). **Coverage:** dedicated reflections from 193847, 829314, 849201, 773921 (Gemini/Claude), the 526815 first-pass, and the 471203 Section-III adversarial batch; substrate attribution inferred from voice where not explicit.

1. Candidate Brief prose / pre-prose.

- The “mathematics of panic” framing for the effects spiral is already adopted into #cooperative-adversarial-intro; it was independently the standout phrase across substrates and is the right Brief anchor for the corollary (Gemini, AUDIT-WORKING-193847). Companion plain gloss: destabilization is “an information-theoretic contest to overflow the opponent's epistemic buffer” — not merely a kinetic tempo race (Gemini, AUDIT-WORKING-829314).
- The product-not-speed correction, stated for a Brief: getting inside the loop “isn't being faster — it's generating disturbance faster than the opponent's absolute structural capacity to absorb it, $\gamma_A \mathcal{T}_A > \Delta \rho_B^*$ ” (Gemini, AUDIT-WORKING-829314; Claude, AUDIT-WORKING-773921).

2. Candidate Discussion.

- **The four levers of destabilization (strongest framing here).** The threshold $\gamma_A \mathcal{T}_A > \alpha_B R_B - \rho_{B,\text{base}}$ exposes four distinct attacker levers, worth naming explicitly in Discussion: (1) increase $\mathcal{T}_A$ — move faster (the brute-force reading); (2) increase γ_A — couple more effectively, act on the variables B cares about most (“moving fast in an irrelevant domain does nothing”); (3) decrease R_B — force B into a narrow, brittle model (the mathematical purpose of deception, feints, conditioning B to predictable patterns); (4) increase $\rho_{B,\text{base}}$ — attack when B 's reserve is already depleted by external chaos or internal reorganization (Gemini, AUDIT-WORKING-829314). The body currently surfaces (1)–(2) via the $\gamma_A \rightarrow 0$ counterpoint; the four-lever decomposition is a candidate Discussion expansion. (*Early-conflation texture: framed as derived strategy doctrine; the levers are real readings of the inequality but the “deception decreases R_B ” link is interpretive.*)
- **Stable-attrition vs structural-collapse threshold.** A 1% tempo edge yields a slightly smaller steady-state mismatch but *both agents survive* — “a stable, predictable war of attrition.” To actually destabilize, A must push ρ_B past B 's model-class capacity R_B entirely. This sharpens the destabilization-vs-steady-state-ratio Discussion paragraph by drawing the line at where attrition becomes collapse (Gemini, AUDIT-WORKING-829314).

3. Follow-up items.

- **Effects-spiral promotion — mechanism candidates (already tracked).** A break-protocol adversarial reading offered concrete candidate mechanisms for the $\gamma_A(\|\delta_B\|)$ functional form whose absence keeps the spiral discussion-grade: (i) erratic actions degrade coupling-channel structure; (ii) degrading M_t produces wider action variance; (iii) other (Claude, AUDIT-WORKING-471203, §III-adversarial). These feed the already-tracked Tier-I spike (spikes/PROPOSED.md, “Effects-spiral eigenvalue condition”) — *not* a new off-ramp, but a useful seed list for whoever runs it. The same reader also proposed a 4th #disc-identifiability-floor instance via Fano’s inequality on observer-side prediction, anchoring H_b as a meta-segment quantity (routes to #der-agent-opacity / the identifiability-floor cluster, not here).

4. Readers often ask / wonder.

- **How does an agent deliberately increase its own opacity H_b without destroying its own ability to act coherently?** (e.g. randomization policies) — a natural defensive question the opacity Discussion invites but does not answer (Gemini, AUDIT-WORKING-849201). Likely belongs with #der-agent-opacity, but the want is generated here.
- **How does an agent defend against a Regime-I-with-adversarial-content (epistemic-poisoning) attack when the incoming data passes all validity checks?** Raised here off the “agent opacity and coupling effectiveness” paragraph; the full treatment is in #der-interaction-channel-classification Discussion, but the question lands on a reader of this segment first (Gemini, AUDIT-WORKING-193847).

5. Candidate figures.

- **Reserve-gauge with side-loop.** A reserve gauge: baseline disturbance consumes part of B ’s reserve $\Delta\rho_B^*$, adversarial coupled tempo $\gamma_A\mathcal{T}_A$ consumes the rest; crossing the threshold exits the invariant region. A side feedback loop shows the effects spiral as an *extra assumption* attached to the base threshold theorem, not part of it — usefully encodes the discussion-grade status visually (Claude, AUDIT-WORKING-526815 first-pass).

Belongs elsewhere.

- **Crèche / kill-switch infrastructure reading of the spiral.** Once the spiral begins ($\dot{V}_B > 0$), the math says B cannot save itself with its own tempo — intervention must come from the infrastructure layer: detect spiral onset and instantly sever the adversarial coupling ($\gamma_A \rightarrow 0$), e.g. take the agent offline or revert it to a safe crèche state (Gemini, AUDIT-WORKING-193847). Aspirational 04-eli-core/ infrastructure reach (resilience = a spiral-breaking mechanism, not just high α/R); develops the consciousness-infrastructure program, not this segment.
- **Monopoly-formation / competitive-investment reading.** The compounding tempo edge is offered as “the thermodynamic reason monopolies form so quickly in tech” — once one firm gets a tempo edge, the squared feedback widens the gap until the slower firm crosses its destabilization threshold and structurally collapses (Gemini, AUDIT-WORKING-829314). This is really a #result-adversarial-tempo-advantage instantiation (squared scaling), filed there; noted here because it co-occurred with the destabilization-threshold reading.

15.3 Derived: Interaction-Channel Classification (Recipient-Side)

The same signal from agent A lands on recipient B as one of four qualitatively different things — informative update, magnitude-shock, structural-shock, or ambient noise — determined by three independent boundary conditions stated entirely in B ’s existing AAT quantities. The emitter-side collapse of this variation into a scalar $\gamma_A\mathcal{T}_A$ loses information that is load-bearing: the recipient’s repair path depends on which regime the event falls into, and “more tempo” vs “different model class” address structurally different failure modes. This is the **anti-collapse** discipline (5.7) at the composition layer — the scalar collapse hides exactly the regime distinction that routes to opposite repairs.

Two event-level quantities enter the classification and *must not be conflated*: the **magnitude** of the event in B ’s observation space (how large a perturbation it produces in mismatch on arrival), and the **information content** of the event

conditional on B 's prior. A large-magnitude already-predicted event has large magnitude but small information content; a tiny-magnitude structurally novel event has small magnitude but large information content. Three independent boundaries — sector-region, model-class, observability-floor — partition events into four regimes. **Regime I (informative update)**: within sector, representable in the model class, above the observability floor — the good case. **Regime II-a (magnitude shock)**: event magnitude exceeds the sector region — too large for the correction mechanism; the repair is more bandwidth or tempo. **Regime II-b (structural shock)**: within sector but information content exceeds the model class — the class lacks representational capacity; the repair is *structural adaptation* (a different model class), and more tempo does not help. **Regime III (ambient erosion)**: below the observability floor — each event individually small enough to be processed but the cognitive overhead accumulates; the repair is *infrastructure-level filtering* before signals reach the agent.

The central methodological move: **the emitter sees a scalar; the recipient sees a regime**. Most multi-agent frameworks treat coupling as a single coefficient (the emitter-side view); the classification names what that coefficient hides — events of identical scalar magnitude can sit in any of the four regimes, with structurally different repairs. The Regime-III ambient-erosion case is the framework's structural reason that DDoS-of-low-priority-alerts is a real attack rather than just inconvenience: not because individual events are significant but because processing them consumes tempo that should go to actual work. For LLM-agent infrastructure this maps directly — push Regime-III filtering *below* the model level into wrapper or system infrastructure rather than processing every event at the model layer. The recipient-side decomposition also surfaces an adversarial move the scalar emitter formulation cannot express, treated in Discussion: *Regime-I-with-adversarial-content* — exploiting B 's openness to informative updates by injecting sign-chosen misinformation on the log-odds signal.

Formal Expression.

Setup and Notation. Two purposeful agents A and B coupled through a shared environment. A 's praxis produces an event e_τ^A that enters B 's observation channel. On B 's side the event is processed by the standard AAT machinery: h_B maps the A -induced environment state to observation o_τ^B (1.3); mismatch is $\delta_\tau^B = o_\tau^B - \hat{o}_\tau^B$; update absorbs δ_τ^B with gain $\eta_B^* = U_{M,B}/(U_{M,B} + U_{o,B})$ (3.6).

Two event-level quantities enter the classification and must not be conflated:

- $\|e_\tau^A\|_B$ — the **magnitude** of the event in B 's observation space (how large a perturbation it produces in δ_τ^B on arrival).
- $\mathcal{I}(e_\tau^A)$ — the **information content** of the event conditional on B 's prior, formally $I(e_\tau^A; \Omega \mid M_{B,\tau-})$ per NOTATION.md's event-information quantity. A large-magnitude already-predicted event has large $\|e\|$ but small $\mathcal{I}$; a tiny-magnitude structurally novel event has small $\|e\|$ but large $\mathcal{I}$.

Let $\mathcal{F}(\mathcal{M}_B)$ denote B 's model-class fitness (2.4), and $\mathcal{I}_{\max}(\mathcal{M}_B)$ the maximum per-event information content representable within the class (see Working Notes for the cleaner sufficient-statistics-span formulation).

Classification Boundaries.

Event e_τ^A arriving at B falls into one of four regimes, determined by three independent boundary conditions:

Definition (regime-boundaries).

Regime I (Informative update) when all three hold:

$$(I-a) \quad \|e_\tau^A\|_B \leq R_B \quad (\text{within sector-condition region}) \quad (15.12)$$

$$(I-b) \quad \mathcal{I}(e_\tau^A) \leq \mathcal{F}(\mathcal{M}_B) \cdot \mathcal{I}_{\max}(\mathcal{M}_B) \quad (\text{representable within model class}) \quad (15.13)$$

$$(I-c) \quad \mathcal{J}(e_\tau^A) \cdot \nu^{(k)} \geq U_{o,B}^{(k)} \cdot c_{\text{floor}} \quad (\text{above observability floor}) \quad (15.14)$$

where k is the arrival channel, $\nu^{(k)}$ its event rate, and c_{floor} a detection-theory constant controlling the false-alarm tolerance.

Regime II-a (Magnitude-shock destabilization) when (I-a) fails:

$$\|e_\tau^A\|_B > R_B \quad (15.15)$$

The event exits B 's sector-condition region on arrival. B 's correction function does not point inward strongly enough to discharge the mismatch before the next event; under sustained rate $\nu > rsim\alpha_B$, destabilization proceeds per 15.2.

Regime II-b (Structural-shock destabilization) when (I-a) holds but (I-b) fails:

$$\|e_\tau^A\|_B \leq R_B, \quad \mathcal{J}(e_\tau^A) > \mathcal{F}(\mathcal{M}_B) \cdot \mathcal{J}_{\max}(\mathcal{M}_B) \quad (15.16)$$

The event's information content exceeds what B 's model class can represent. By 4.8, parametric update within $\mathcal{M}_B$ cannot close the mismatch; residuals retain systematic structure. Repair requires structural adaptation (a different model class), not more bandwidth.

Regime III (Ambient noise / slow erosion) when (I-a) and (I-b) hold but (I-c) fails:

$$\|e_\tau^A\|_B \leq R_B, \quad \mathcal{J}(e_\tau^A) \leq \mathcal{F}(\mathcal{M}_B) \cdot \mathcal{J}_{\max}(\mathcal{M}_B), \quad \mathcal{J}(e_\tau^A) \cdot \nu^{(k)} < U_{o,B}^{(k)} \cdot c_{\text{floor}} \quad (15.17)$$

The event is representable and within capacity but its information content sits below the observability floor. It contributes to δ_B 's variance (enters Model S as part of $\sigma_{w,B}^2$) without triggering a usable update; B 's adaptive reserve $\Delta\rho_B^*$ slowly drains.

Three Independent Boundaries. The three boundary conditions are structurally independent, each stated in quantities AAT already carries:

Table 15.1

<i>Boundary</i>	<i>AAT quantities</i>	<i>Failure mode</i>
(I-a) / (II-a): sector-region	$\ e\ _B, R_B$ (from 2.4 / G)	<i>magnitude</i> — more capacity cures
(I-b) / (II-b): model-class	$\mathcal{J}(e), \mathcal{F}(\mathcal{M}_B), \mathcal{J}_{\max}(\mathcal{M}_B)$	<i>class</i> — structural adaptation cures
(I-c) / (III): observability	$\mathcal{J}(e), \nu_B^{(k)}, U_{o,B}^{(k)}$ (from 16.4)	<i>rate</i> — lower observation noise or higher event rate cures

No new ad-hoc thresholds are introduced. $\mathcal{J}_{\max}(\mathcal{M}_B)$ is the only new symbol; see Working Notes for its cleaner sufficient-statistics-span formulation.

Regime-Typed Disturbance Decomposition.

Under a stream of events $\{e_\tau^A\}$, B 's **regime-typed effective disturbance** rate decomposes into regime-typed contributions:

Derived (regime-typed-rho-eff, from regime-boundaries + sector-persistence-template).

$$\rho_B^{\text{eff}} = \underbrace{\sum_{e \in \text{II-a}} \|e\|_B \cdot \nu_e}_{\text{magnitude disturbance}} + \underbrace{\text{floor}(\mathcal{M}_B) \cdot \sum_{e \in \text{II-b}} \nu_e}_{\text{structural mismatch floor}} + \underbrace{\sum_{e \in \text{III}} \sigma_e^2 \cdot \nu_e}_{\text{ambient variance}} - \underbrace{\sum_{e \in \text{I}} \iota_B(e) \mathcal{J}(e) \cdot \nu_e}_{\text{informative correction}} \quad (15.18)$$

The Regime-I term is **negative**: informative events reduce B 's effective disturbance rate, not increase it. This generalizes 15.1's cooperative-action term $-\gamma^{\text{coop}} \mathcal{J}$: a cooperative event is precisely a Regime-I event from an aligned emitter, and the sign flip in the emitter-side decomposition falls out of the regime assignment on the recipient side. Adversarial events land in Regimes II-a/II-b; ambient-noise events in Regime III.

The emitter-side formulation $\gamma_A \mathcal{J}_A \rightarrow \rho_B^{\text{eff}}$ compresses the regime-typed sum into a single scalar, losing (i) the sign of cooperative coupling, (ii) the magnitude vs structural distinction in destabilization, and (iii) the observability-floor loss to Regime III.

Structured Derivation — Kalman-over-Kalman.

For a concrete check, take B as a Kalman filter on a scalar linear-Gaussian state with model class $\mathcal{M}_B = \{\theta \in [\theta_{\min}, \theta_{\max}]\}$, process noise q , observation noise r , sector parameter $\alpha_B = \eta_B^* = P_{\text{pred}}/(P_{\text{pred}} + r)$. B 's sector-region radius is $R_B = \sqrt{q/(1 - \theta_{\max}^2)}$ (stationary standard deviation at the class edge).

Derivation (Kalman-over-Kalman, from regime-boundaries + update-gain).

A 's emitted perturbation ξ_A enters B 's innovation as $\delta_\tau^B = \xi_A + \varepsilon_\tau + (\omega_\tau - \hat{\omega}_\tau)$. Four canonical distributions for ξ_A partition the classification:

Case 1 — Small-variance Gaussian $\xi_A \sim \mathcal{N}(0, s^2)$, $s^2 \ll r$ (**expected Regime III**). $\mathcal{J}(\xi_A) \approx s^2/(2r \ln 2)$ nats — small. (I-a) holds ($s \ll R_B$); (I-b) holds (Gaussian-within-Gaussian); (I-c) fails because $\mathcal{J}(\xi_A) < U_{o,B} \cdot c_{\text{floor}}$. Result: Regime III. Derived consequence — the contribution to ρ_B^{eff} is through $\sigma_{w,B}^2$; adaptive reserve drains by $\sum_e \eta_B^{*2} s_e^2 \cdot \nu_e$.

Case 2 — Moderate Gaussian $\xi_A \sim \mathcal{N}(\mu, s^2)$, $\mu \ll R_B$, $s^2 \sim r$ (**expected Regime I**). $\mathcal{J}(\xi_A) = \frac{1}{2} \log(1 + (s^2 + \mu^2)/r)$ — substantial. All three (I-a)/(I-b)/(I-c) hold. Result: Regime I. Standard Kalman update; M_B refines; Regime-I term contributes negatively to ρ_B^{eff} .

Case 3 — Binary kick $\xi_A \in \{\pm\Delta\}$ with $\Delta > R_B$ (**expected Regime II-a**). (I-a) fails by construction. The Kalman update $\hat{x}^+ = \hat{x}^- + \eta^* \Delta$ undershoots by $(1 - \eta^*)\Delta$ per event. If events arrive at rate $\nu > r \text{sim} \alpha_B$, lag accumulates and $\alpha_B R_B > \rho_B^{\text{eff}}$ is violated. Result: Regime II-a — destabilization per 15.2. Notice: the signal is *within* the model class (Gaussian handles $\pm\Delta$ mathematically), but correction cannot discharge it fast enough.

Case 4 — Heavy-tailed ξ_A with $\mathbb{E}[\xi_A^2] \sim r$ but kurtosis $\kappa \rightarrow \infty$ (expected Regime II-b). Mean contribution is fine; the problem is the distribution shape. The Kalman filter — Gaussian-optimal — mis-gains: too aggressive for small events, too conservative for genuine large ones. The per-event KL gap $D_{\text{KL}}(P_{\text{true}} \| P_{\mathcal{M}_B}) > 0$ for any heavy-tailed P_{true} against Gaussian. By 2.4, $\mathcal{F}(\mathcal{M}_B) < 1 - \varepsilon$ with ε lower-bounded by the KL gap; by 4.8, no parametric update within $\mathcal{M}_B$ closes the mismatch. Result: Regime II-b — residuals retain non-Gaussian structure (visible in kurtosis tests); repair requires expanding the model class (e.g., Student- t observation model), not more Kalman tuning.

Each case lands where the classification predicts. The derivation transfers to any recipient architecture in which the underlying AAT quantities are well-defined — this is the scope inherited from G + 2.4 + 3.8.

Recovery of Emitter-Side Results.

Each existing emitter-side result is a restriction of the four-regime decomposition:

Derived (emitter-side-recovery).

- **15.2** is Regime II-a integrated over a tempo-proportional event stream. The magnitude-shock sub-regime corresponds directly; the structural-shock II-b subcase is implicit in that segment, collapsed into “adaptive reserve exceeded” but here made explicit.
- **15.4** — superlinear tempo scaling follows from the sector-persistence template’s $1/\alpha$ (Model D) vs $1/\sqrt{\alpha}$ (Model S) applied to Regime II events. The b -exponent drops toward zero in the high- $U_{o,B}$ limit because the fraction of A ’s events landing in Regime II drops (more fall into Regime III).
- **16.4** is the recipient-side expression of boundary (I-c): high $U_{o,B}$ pushes events into Regime III where they add to variance without contributing to destabilization.
- **12.3** corresponds to asymmetric classification: host’s signals to endosymbiont contain high- $\mathcal{J}$ structure the endosymbiont’s class cannot initially represent (Regime II-b for the endosymbiont, forcing structural adaptation toward the host’s class); endosymbiont’s signals to the host land in Regime I (host absorbs endosymbiont’s accumulated structure). Consolidation is the fixed-point where both streams are Regime I.
- **Cooperative signaling** (in 15.1) — Regime-I events from aligned emitters contribute negatively to ρ_B^{eff} via the cooperative-action term. The communication-tempo $\nu_{ji}^{\text{comm}} \cdot \eta_{ji}^*$ is the rate of Regime-I events times the recipient’s informative gain.

What Is Derived vs. What Is Chosen.

Table 15.2

<i>Property</i>	<i>Source</i>	<i>Strength</i>
Three independent boundaries (sector-region / model-class / observability)	Import from 2.4 + G + 16.4	Formulation choice (three-way partition; coarser / finer alternatives possible)
Four-regime partition (I / II-a / II-b / III)	The three boundaries yield four boundary-state combinations; only four are non-degenerate	Derived from the boundary structure
(I-a) / (II-a) boundary at R_B	G’s sector-region radius	Derived
(I-b) / (II-b) boundary at $\mathcal{F}(\mathcal{M}_B) \cdot \mathcal{J}_{\max}$	2.4 + class-capacity normalization	Formulation (the $\mathcal{J}_{\max}$ normalization is heuristic; sufficient-statistics-span form is cleaner — see Working Notes)
(I-c) / (III) boundary at $U_{o,B} \cdot c_{\text{floor}}$	16.4 + detection-theory threshold	Formulation (c_{floor} is a detection-power parameter)
Regime-typed ρ_B^{eff} decomposition with negative Regime-I term	Aggregation over event stream using regime classification	Derived (the sign of the Regime-I term is structural, not a choice)
Kalman-over-Kalman four-case derivation	Direct application of Kalman gain + sector + KL-gap + SNR analyses	Proved (for the stated case)
Recovery of 15.2, 12.3, 16.4, 15.1 as restrictions	Each emitter-side result is exhibited as a per-regime special case	Derived
Non-Gaussian Case 4 derivation (heavy-tailed $\rightarrow$ II-b via KL gap)	Informal argument grounded in robust-filtering literature (Huber, Masreliez)	Discussion-grade (rigorous version requires per-family KL computation)
Class 2 (Partial) approximation with $\kappa_{\text{processing}}$ degradation	Transfer from Class 1 (Separated) with goal-blind-update failure	Formulation (qualitative); exact form requires spelling out the goal-contamination coupling

Epistemic Status.

Conditional. Max attainable: *derived* for the classification structure in Class 1 (Separated) architectures; *exact* for the Kalman-over-Kalman worked case in sub-scope α ; *robust qualitative* for general sub-scope β recipients.

The three boundaries are structurally independent and each is stated in an existing AAT quantity, so the four-regime partition is not ad-hoc; it is forced by the structure of the quantities already in the theory. Within Class 1 (Separated; goal-blind epistemic update per 6.3) and sub-scope α recipients (Kalman / conjugate-Bayesian / exponential-family / strongly-convex-gradient / linear-PD), the Kalman-over-Kalman derivation transfers with A2' derived per 4.3, and each case yields the predicted regime by direct substitution. In sub-scope β (PID / rule-based / human-judgment), the classification's form transfers but boundary (I-a) requires per-instantiation verification of the sector bound, inheriting from G 's sub-scope β caveat.

Scope limits. Class 3 (Coupled) recipients are out of formal scope: the coupled epistemic-purposeful update violates 6.3, so the regime assignment cannot be cleanly computed against M_B alone. This is 03-11m-core/ territory. Class 2 (Partial) recipients inherit with degradation proportional to $\kappa_{\text{processing}}$: the Regime-I update may be goal-contaminated, and Regime II-b may be misdiagnosed as II-a or III under goal-blind-update failure. Per-event classification is singular-trajectory-indexed by 4.10; aggregation over event streams is deferred to G .

What the classification does not claim. (i) It does not determine *semantic content*: knowing an event is Regime I tells you B will update, not *what* B will believe afterward — that lives in 9.6's default signal function. (ii) It does not make the four regimes sharp in practice: real recipients' sector regions and class capacities have estimation uncertainty, so borderline events oscillate across boundaries under parameter noise. (iii) It does not replace G 's temporal aggregation; it provides the per-event regime label that feeds into the aggregate disturbance bound.

Discussion.

Complement, not replacement, of the emitter-side segments. 15.2, 15.4, 16.4, 12.3, and 15.1 each describe what happens when A 's praxis couples into B . They collapse the variation into a scalar increment on B 's effective disturbance rate. The recipient-side classification decomposes that increment into regime-typed contributions with (importantly) a **signed** Regime-I component. The five emitter-side segments are recovered as restrictions — the classification provides the pattern-recognizer to match the emitter-side optimizer.

Why II-a vs II-b matters. Both destabilization regimes manifest as “adaptive reserve exceeded” in the emitter-side view. But the repairs are different: magnitude-shock (II-a) calls for more bandwidth — larger R_B , higher α_B , faster tempo; structural-shock (II-b) calls for a different model class — an expansion, grafting, or compression in the sense of 4.8. An organization running rapid-response incident drills against a bureaucratic-process adversary faces II-a and responds with more capacity; the same organization hit by a technology-regime change faces II-b and responds with restructuring. The two produce similar pain signals but admit opposite cures. Collapsing them confuses diagnosis.

Pairing with #adversarial-edge-targeting (Part III GAP). The classification provides the recipient-side pattern-recognizer; the emitter-side question — which of B 's strategy edges are most valuable for A to attack — provides the adversary-optimization complement. The two compose: the emitter chooses which edge to target; the classification determines what happens at that edge. A particularly sharp move the classification makes visible but the emitter-side formulation hides: the **Regime-I-with-adversarial-content attack** — an emitter who can control the content y_G fed into B 's default log-odds signal function (per 9.6) can choose the sign of $(y_G - \hat{R}_2)$ to push λ_k in the direction that degrades Σ_B most. This is not a destabilizing attack (which lands in Regime II); it is an *informational* attack that exploits B 's openness to Regime-I updates to inject misinformation. The emitter-side formulation, which represents coupling as a scalar γ_A , cannot express the sign lever; the recipient-side decomposition does. The value is maximized at edges with large plan-sensitivity J_k and moderate credence p_k — exactly the edges load-bearing for Σ_B . The full arg-max — these two factors plus edge identifiability ι_k , observability σ_k^B , and the opacity targeting-fidelity factor $(1 - H_b^{B|A}/H_b^{\max})$ — is displayed and grounded in 16.2 §”The adversarial-edge-targeting arg-max” (the emitter-side optimizer this recipient-side classifier pairs with).

Agent opacity and the emitter-side regime distribution. Hafez et al. (2026)’s agent opacity $H_b^B = H(S, A \mid S')$ is primarily *emitter-side*: it gates A ’s ability to model which of B ’s regimes a given signal will land in. As $H_b^B \rightarrow 0$ the emitter can target specific regimes adversarially; as $H_b^B \rightarrow \infty$ the emitter’s signals distribute across regimes and the targeting advantage collapses. Formally, $\gamma_A^{\text{effective}} = \gamma_A^{\text{max}} \cdot f(H_b^B)$ with f monotonically decreasing. This is the dual of 16.4’s story: where observation-gates-advantage says the emitter’s advantage degrades when its observations of the *environment* are noisy, opacity says the advantage degrades when its observations of B are noisy. Cooperative signaling is the mirror: low H_b^B (legibility) enables aligned emitters to target signals into B ’s Regime I. This is 14.3’s recipient-side translation: shared objectives give A a better model of B ’s prior and model class, which lets A produce signals that land in B ’s informative regime.

Organizational reception intuition. The four regimes correspond to a familiar organizational diagnostic: (Regime I) *heard* — the organization updated on the relevant edges of its Σ_{org} ; (II-a) *shattered by shock* — signal was too large/fast, the organization’s correction bandwidth was exceeded; (II-b) *couldn’t hear* — signal required reorganization beyond the organization’s structural capacity, matching Cohen & Levinthal’s (1990) absorptive-capacity framing; (III) *absorbed as fatigue* — signal drained $\Delta\rho^*$ without producing structural change, the “death by a thousand memos” pattern. The cross-domain instantiation is cleanest where $\mathcal{F}(\mathcal{M}_B)$ maps to absorptive capacity and $U_{o,B}$ to communication-channel quality.

Future meta-segment candidate. The recipient-side classification is the *signal-reception* complement to G ’s *persistence-under-disturbance* structure. If a future pattern emerges elsewhere (e.g., consolidation dynamics classifying inputs, or a more general theory of inter-agent coupling) a meta-segment #signal-reception-pattern could name the recipient-classification shape. Premature at this stage; log as speculation.

Working Notes.

- **Migration note (2026-05-09 GUC rename):** Class 2 $\leftrightarrow$ Class 3 swap. Pre-2026-05-09: Class 2 = fully merged, Class 3 = partially modular. Post: Class 2 = Partial, Class 3 = Coupled. Scope-limits paragraph updated (Class 2 fully-merged $\rightarrow$ Class 3 Coupled; Class 3 partially-modular $\rightarrow$ Class 2 Partial). What-Is-Derived table updated (Class 3 approximation $\rightarrow$ Class 2 Partial approximation). The four-regime classification (I / II-a / II-b / III) is recipient-side signal-processing vocabulary — unaffected by the GUC rename. Removed at candidate stage per FORMAT.md Gate 4.
- $J_{\text{max}}(\mathcal{M}_B)$ — **replace with sufficient-statistics-span.** The boundary (I-b) uses $\mathcal{F}(\mathcal{M}_B) \cdot J_{\text{max}}(\mathcal{M}_B)$ as the class’s representational ceiling on per-event information. The cleaner formulation: (I-b) holds iff the event’s sufficient statistics for prediction lie in the span of $\mathcal{M}_B$ ’s sufficient statistics. For parametric families (exponential families, Gaussian) the explicit form is routine; for non-parametric classes it requires the projection-to-class formalism of 2.2. Worth refining at the next candidacy review.
- **Four regimes vs finer splits.** The partition is the *coarsest* useful split where each regime has a distinct AAT-machinery response. Regime I could split into within-observability / within-calibration (point update vs structural refinement) as a fifth-regime refinement; Regime II-a could split by transient vs sustained dynamics. These duplicate the structural-adaptation hierarchy and are not pursued here.
- **Heavy-tailed derivation rigorization.** The Case-4 argument is informal; a rigorous version would compute the effective Kalman gain under heavy-tailed observation misspecification and show residuals retain signal. This is the robust-filtering territory (Huber, Masreliez) and is standard; the classification does not depend on a specific non-Gaussian family, just on the KL gap being nonzero.
- **Regime II-b as candidate 5.1 Instance 3.** Under a sustained Regime II-b stream, B ’s ρ_B^{eff} degrades at a rate lower-bounded by the KL gap between $\mathcal{M}_B$ and the true event distribution — a quantitative misspecification-cost floor. This is adjacent to 5.1’s “Misspecification-cost quantification” open extension. Worth a focused spike to formalize.
- **Connection to active inference.** Regime II-b corresponds to high-surprise events that exceed the agent’s generative model — what AI calls “novel generative context.” The classification avoids AI’s variational-free-energy-as-master commitment (per the

spike-routing cycle, CHANGELOG 2026-05-17); it uses only the information-theoretic content of $J(e)$, not the variational machinery.

- **Next spike candidates.** (1) Formalize the $f(H_b^B)$ emitter-side-effect function — tighten the qualitative $\gamma_A \propto f(H_b^B)$ into a derived form (tracked: TODO.md “Queued spike work”; cross-referenced in spikes/PROPOSED.md Reserved/owned-elsewhere). (2) Formalize Regime II-b misspecification-cost as a structural result (5.1 Instance 3 candidate) (tracked: TODO.md identifiability-floor queued-spike cluster; same cross-reference). (3) **Resolved — no longer a candidate.** The #adversarial-edge-targeting arg-max *is* formally solved in canon: 16.2 derives it as the 16-cell emitter-recipient composition (a Derived arg-max; the previously-reserved Section-III GAP is filled), with *this* segment supplying the recipient-side four-regime classification that pairs with the emitter side. Scope per 16.2 Epistemic Status: closed-form under sub-scope- α coupling; general non-convex coupling is per-case optimization, a noted scope caveat, not an open derivation.

Incidental audit gold (lift 2026-05-31). Cross-audit “wandering thoughts” / §14-ideation harvested from the de-novo auditors’ working dirs, deduplicated across substrates and attributed by substrate + audit. Orthogonal pedagogical / framing / figure / naming material, kept separate from the certified theory-fix findings (the 526815 F203–F213 stream — event-information realized-vs-expected notation, the ρ_B^{eff} common-unit ledger, the obs-gated-tempo-advantage undeclared dependency, the Case-I nats/bits unit slip — is certified-track and routed there; see also the Off-ramp block below for one structural-precision item). **Coverage:** dedicated reflections from 193847, 829314, 849201 (Gemini), the 526815 first-pass, and the 451729 batch-14; substrate attribution inferred from voice where not explicit.

1. Candidate Brief prose / pre-prose.

- The “emitter sees a scalar; the recipient sees a regime” line is already the segment’s central methodological hook; across substrates it reads as the standout. A complementary plain gloss for a Brief: the four regimes are an organizational diagnostic — “*heard* (Regime I), *shattered by shock* (II-a), *couldn’t hear* (II-b), *absorbed as fatigue* (III)” (the organizational-reception mapping in the segment’s own Discussion, independently praised by Gemini, AUDIT-WORKING-193847).

2. Candidate Discussion.

- **Phishing / social-engineering isomorphism for the Regime-I-with-adversarial-content attack (strongest framing here).** The segment names the sign-lever attack abstractly; three substrates independently reached for the same concrete instance and it is sharp enough to seed a Discussion example. A brute-force attack (DDoS, bombardment) is Regime II-a — it hits the firewall R_B and is rejected, loud and obvious. A phishing email or data-poisoning sample is Regime I — to the learning machinery f_M the data looks *great*: inside the sector region, representable in $\mathcal{M}_B$, above the observability floor, so the gain η^* happily ingests it. Because the content is adversarially chosen to maximize strategic error, “you are efficiently optimizing your way into a trap — the faster you learn, the faster you die.” This is why social engineering beats brute-force hacking: it hijacks the agent’s own cooperative-update rules to breach from the inside (Gemini, AUDIT-WORKING-829314; Gemini, AUDIT-WORKING-849201; Gemini, AUDIT-WORKING-193847). (*Early-conflation texture: presented as a derived security result; it is a vivid instantiation of the existing discussion-grade sign-lever claim, not a new derivation.*)
- **The opacity defense against Regime-I poisoning.** The named defense against the sign-lever attack is agent opacity: hide internal model structure from the adversary so they cannot craft the perfect Regime-I poison (the adversary needs a model of B ’s prior and class to choose the damaging sign). Connects the recipient-side attack directly to #der-agent-opacity’s H_b (Gemini, AUDIT-WORKING-829314).

3. Follow-up items.

- **Diagnostic-conflation watch as a candidate explicit warning.** Multiple substrates flagged the II-a/II-b conflation as the load-bearing reader trap: both manifest as “adaptive reserve exceeded,” but a downstream claim that assumes high mismatch is always cured by more tempo is wrong (II-b needs a new model class). The segment’s Discussion makes this point; the auditors’ convergence suggests stating it as an explicit boxed warning where the downstream-repair logic is introduced (Gemini, AUDIT-WORKING-193847; Gemini, AUDIT-WORKING-829314; Gemini, AUDIT-WORKING-849201).

- **Sufficient-statistics-span formalization.** A reader specifically wanted the cleaner $J_{\max}$ replacement (sufficient-statistics span) developed — already flagged in this segment’s own Working Notes; the fresh-reader want confirms it is worth prioritizing at the next candidacy review (Gemini, AUDIT-WORKING-849201).

4. Readers often ask / wonder.

- **Can the theory predict *which* regime (II-a vs II-b) an adversary will choose to induce?** A natural strategic question once the two destabilization sub-regimes are distinguished — a destabilized agent might still be in the right model class with saturated parameters (II-a) vs needing a paradigm shift (II-b) (Claude, AUDIT-WORKING-773921, raised at #der-adversarial-destabilization).
- **How is Regime-I-with-adversarial-content prevented without infinite U_o (paranoia)?** If the poison passes all validity checks, what stops absorption short of refusing to learn? (Gemini, AUDIT-WORKING-193847; Gemini, AUDIT-WORKING-829314).

5. Candidate figures.

- **Precedence-flow classifier.** A precedence flowchart — first check sector magnitude (R_B), then representability ($\mathcal{F} \cdot J_{\max}$), then observability floor — terminating in the four named regimes, with each leaf annotated by its distinct repair (more bandwidth / new model class / better filtering / ordinary update). The flowchart was independently judged “clearer than the quantitative aggregation” (Claude, AUDIT-WORKING-526815 first-pass). (*Note: drawing it as a precedence flow also makes visible the structural-precision point in the Off-ramp block below.*)

Belongs elsewhere.

- **Regime-III “firewall” for logogenic-agent infrastructure.** Under a Regime-III flood (DDoS-of-low-priority-alerts), the agent’s tempo is consumed parsing noise it must still *process* to discover is below-floor. The infrastructure prescription: aggregate Regime-III noise over long windows ($K_c \gg 1$) at the wrapper/system layer and only escalate to the core engine if the aggregate crosses the observability threshold — “otherwise the intelligence is paralyzed by the ambient hum of its own sensors” (Gemini, AUDIT-WORKING-193847). The segment already names infrastructure-level filtering as the Regime-III repair; this develops the 03-11m-core/ PROPRIUM-ASM / attention-management implementation, not this segment’s formal content.

Off-ramp (NOT gold — flagged for the certified track).

- **Structural-precision point on “the four regimes are forced.”** The segment states the four-regime partition is “forced by the structure of the quantities already in the theory” (Epistemic Status; “Three Independent Boundaries”). 526815’s first-pass (F205/F187) and a fresh reading both observe that three boolean boundary tests yield more than four boolean combinations; the four named regimes follow from a *precedence convention* (fail magnitude $\Rightarrow$ II-a; else fail class $\Rightarrow$ II-b; else fail observability $\Rightarrow$ III; else I), and two of the tests both depend on $\mathcal{I}(e)$ so the boundaries are *orthogonal diagnostic criteria* rather than statistically independent (526815 F187). This is not framing gold — it is a precision point about a load-bearing structural claim, and the honest repair is to state the precedence/coarsening explicitly and soften “forced by independent boundaries” to “the coarsest useful partition under this precedence” (the Working Notes “Four regimes vs finer splits” bullet already gestures at coarseness). Routed to the certified-findings track for adjudication; recorded here because the claim is structural, not merely a convention.

15.4 Result: Adversarial Tempo Advantage

Under adversarial coupling where one agent’s actions contribute to the other’s disturbance rate, the steady-state mismatch ratio scales superlinearly with the tempo ratio. The construction applies the sector-persistence template’s steady-state formula to both agents and ratios them; in the *coupling-dominant limit* with symmetric coupling, the result reduces to a clean power law.

Under Model D (deterministic drift coupling), the exponent is $b = 2$ — the squared tempo advantage. A 2:1 tempo ratio yields a 4:1 mismatch ratio; 3:1 yields 9:1; 10:1 yields 100:1. The faster agent both corrects its own mismatch faster *and* generates disturbance for the opponent faster — the two effects compound multiplicatively rather than add. **Under Model S (stochastic noise coupling) the exponent drops to $b = 3/2$** , inheriting the $1/\sqrt{\alpha}$ steady-state scaling from A Prop A.1S. **Under a non-coupling-dominant regime** (where base disturbance dominates coupling-induced disturbance) the exponent approaches 1 (Model D) or $1/2$ (Model S). The regime where the squared scaling is sharpest is also where it matters most.

The central methodological framing is that this is **the formal analog of Boyd’s “getting inside the opponent’s OODA loop”** — military intuition becomes physics. The math itself is not mysterious; it is the steady-state formula applied to both sides and ratioed. The consequence is sharp: at 100:1 mismatch ratio the slower agent is almost certainly past its operating reserve R , which means **structural collapse rather than degraded performance**. Speed advantage in a coupled regime *compounds*, and the consequence is qualitatively different from what additive intuition predicts. The result is conditional on stated scope: the linear-correction case ($\alpha = \mathcal{T}$); the coupling-dominant limit; and the disturbance model. Paired with the scope-honest counterpoint from 15.2 — **speed without coupling is useless**: infinite tempo with zero coupling effectiveness destabilizes no one. Speed and coupling are *jointly required*; the squared scaling only kicks in when both are present in the coupling-dominant regime.

Formal Expression.

Setup. Two agents A, B with adaptive tempos $\mathcal{T}_A, \mathcal{T}_B$, each instantiating G with linear correction ($\alpha = \mathcal{T}$). The adversarial coupling of 15.2 enters each agent’s effective disturbance:

Derived (adversarial-tempo-advantage, from sector-persistence-template + adversarial-destabilization coupling model).

$$\rho_A^{\text{eff}} = \rho_{\text{base}} + \gamma_B \cdot \mathcal{T}_B, \quad \rho_B^{\text{eff}} = \rho_{\text{base}} + \gamma_A \cdot \mathcal{T}_A \quad (15.19)$$

with $\gamma_A, \gamma_B > 0$ coupling effectivenesses and ρ_{base} the shared background rate (asymmetric ρ_{base} generalizes straightforwardly).

Model D: Deterministic coupling, $b = 2$.

The template’s Model D conclusion $\|\delta\|_{ss} = \rho^{\text{eff}}/\mathcal{T}$ (linear case, $\alpha = \mathcal{T}$) applied to both agents and ratioed:

Result (adversarial-tempo-advantage, Model D).

$$\frac{\|\delta_B\|_{ss}}{\|\delta_A\|_{ss}} = \frac{(\rho_{\text{base}} + \gamma_A \mathcal{T}_A) \mathcal{T}_A}{(\rho_{\text{base}} + \gamma_B \mathcal{T}_B) \mathcal{T}_B} \quad (15.20)$$

In the coupling-dominant limit ($\gamma\mathcal{T} \gg \rho_{\text{base}}$) with symmetric coupling ($\gamma_A = \gamma_B$):

$$\frac{\|\delta_B\|_{ss}}{\|\delta_A\|_{ss}} \rightarrow \left(\frac{\mathcal{T}_A}{\mathcal{T}_B}\right)^2 \quad (15.21)$$

The exponent is $b = 2$: a **squared** tempo advantage. A 2:1 tempo ratio yields a 4:1 mismatch ratio. The faster agent both (a) corrects its own mismatch faster and (b) generates disturbance for the opponent faster — the two effects compound rather than add. $\square$

Model S: Stochastic coupling, $b = 3/2$.

Under Model S the coupling enters the noise scale: $\sigma_B^{\text{eff}} = \sigma_{\text{base}} + \gamma_A \mathcal{T}_A$. Adversary tempo increases unpredictability, not systematic direction. The template's Model S steady state $\|\delta\|_{\text{rms}} = \sigma/\sqrt{2\mathcal{T}}$ (linear $\alpha = \mathcal{T}$, scalar $n = 1$) applied to both agents:

Derived (stochastic-tempo-advantage, from sector-persistence-template Model S + coupling).

$$\frac{\|\delta_B\|_{\text{rms}}}{\|\delta_A\|_{\text{rms}}} = \frac{(\sigma_{\text{base}} + \gamma_A \mathcal{T}_A)\sqrt{\mathcal{T}_A}}{(\sigma_{\text{base}} + \gamma_B \mathcal{T}_B)\sqrt{\mathcal{T}_B}} \quad (15.22)$$

In the coupling-dominant, symmetric limit:

$$\frac{\|\delta_B\|_{\text{rms}}}{\|\delta_A\|_{\text{rms}}} \rightarrow \left(\frac{\mathcal{T}_A}{\mathcal{T}_B}\right)^{3/2} \quad (15.23)$$

The exponent is $b = 3/2$. $\square$

Why 3/2, not 2. The half-power difference between the template's Model D ($1/\alpha$) and Model S ($1/\sqrt{\alpha}$) scalings propagates through the ratio. Numerator contributes $\mathcal{T}_A^1$ from the coupling; denominator contributes $\mathcal{T}_B^{1/2}$ from noise averaging; combined with the A -side $1/\mathcal{T}_A^{1/2}$ gives $\mathcal{T}_A^{3/2}/\mathcal{T}_B^{3/2}$.

Summary of Regime-Dependent Exponents.

Table 15.3

<i>Regime</i>	<i>Coupling type</i>	<i>Dominance</i>	<i>Exponent b</i>	<i>Source</i>
1	Deterministic drift (Model D)	Coupling-dominant	2	Derived above
2	Stochastic noise (Model S)	Coupling-dominant	3/2	Derived above
3	Either	Non-coupling-dominant	$\rightarrow 1$ (det.) or $\rightarrow 1/2$ (stoch.)	Asymptotic limit

Regime 3 (non-coupling-dominant). When $\rho_{\text{base}} > r\text{simy} \cdot \mathcal{T}$ (or $\sigma_{\text{base}} > r\text{simy} \cdot \mathcal{T}$), the base disturbance dominates and the coupling terms become a perturbation. The mismatch ratio degrades toward $\mathcal{T}_A/\mathcal{T}_B$ (linear, $b = 1$) for Model D, or toward $(\mathcal{T}_A/\mathcal{T}_B)^{1/2}$ for Model S.

The simulation validation across all three regimes is in 16.3.

Epistemic Status.

Both coupling-dominant exponents are *exact* conditional on their respective disturbance models. The squared law ($b = 2$) is *exact* under Model D (deterministic bounded disturbance, GA-2) with coupling-dominant conditions. The 3/2 law ($b = 3/2$) is *exact* under Model S (stochastic disturbance, GA-2S) with coupling-dominant conditions, derived from the $1/\sqrt{\alpha}$ steady-state scaling (Prop A.1S in A). Both derivations are straightforward algebra from the respective steady-state formulas and the coupling model. The coupling model itself is an *assumption* — the same one used in 15.2.

The non-coupling-dominant limits ($b \rightarrow 1$ for Model D, $b \rightarrow 1/2$ for Model S) are derived asymptotically. The smooth transition between regimes is confirmed by simulation (16.3) but the interpolation formula is empirical. The transition between regimes is smooth, not sharp.

Max attainable: exact conditional on the disturbance model and coupling model. The result is as strong as its assumptions; no additional work changes the epistemic status without changing the dynamical model.

Discussion.

Superlinearity is the key result. The naive expectation — twice as fast yields twice the advantage — is wrong under adversarial coupling. The mechanism is that the faster agent both (a) corrects its own mismatch faster and (b) generates disturbance for the opponent faster. These two effects multiply, producing the squared exponent. Speed advantage is not additive; it compounds.

Relationship to 15.2. The steady-state mismatch ratio quantifies how much worse the slower agent does *while both agents persist*. The destabilization threshold (15.2) marks where the slower agent fails entirely — its correction mechanism breaks down. Below the threshold, this segment's mismatch ratio applies. Above it, 15.2's Lyapunov divergence takes over. The two results are complementary: this one gives the score; that one gives the game-ending condition.

Regime dependence is operationally significant. Whether an adversary's tempo increase produces systematic drift (positional maneuvering, API changes, doctrinal initiative) or unpredictable noise (feints, randomized attacks, market volatility) determines the scaling law. The distinction is not academic — $b = 2$ vs. $b = 3/2$ means a 3:1 tempo ratio yields 9:1 vs. 5.2:1 mismatch ratio. The model predicts that consistent, directional pressure is more effective per unit of tempo than unpredictable disruption.

Formal analog of OODA-loop observations. The squared scaling is consistent with Boyd's observation that getting inside the opponent's decision cycle has disproportionate effects. The theory identifies a specific mechanism (multiplicative interaction of correction speed and disturbance generation) and a specific condition (coupling-dominant regime) under which this disproportionality holds. Whether this mechanism is the dominant one in actual adversarial interactions is an empirical question, not a mathematical one.

Working Notes.

- **Channel-independence assumption.** The tempo ratio $\mathcal{T}_A/\mathcal{T}_B$ uses scalar tempo, which inherits the channel-independence assumption from 3.8. When either agent's observation channels are correlated, the additive formula overcounts their tempo, inflating or deflating the ratio and the derived mismatch advantage. The superlinear exponents ($b = 2$, $b = 3/2$) are exact given the scalar tempos; the caveat concerns whether the scalar tempos themselves are accurate.
- **Opacity modulation (16.2).** The coupling effectiveness γ_A is not primitive: it decomposes as $\gamma_A^{\text{adv, effective}} \propto H_b^{A|B}/H_b^{\text{max}}$, so this segment's tempo $\mathcal{T}_A$ enters adversarial dynamics as an *effective* tempo $\mathcal{T}_A^{\text{effective}} = \mathcal{T}_A \cdot H_b^{A|B}/H_b^{\text{max}}$ (Model D). The superlinear ratio then carries a squared bilateral-opacity factor $(H_b^{A|B}/H_b^{B|A})^2$; the exponent b is unchanged (opacity is a prefactor, not an exponent — open whether bilateral opacity reshapes b). A slow-but-opaque adversary can match a fast-but-transparent one at $\mathcal{T}_A H_b^{A|B} \approx \mathcal{T}_B H_b^{B|A}$.

- The analysis treats each agent's tempo as exogenous — $\mathcal{T}_A$ does not change in response to B 's actions and vice versa. A fully coupled analysis where both agents' mismatch states co-evolve simultaneously (joint Lyapunov function over (δ_A, δ_B)) is the open extension. The decoupled result is a worst-case bound for the slower agent: in practice, the faster agent may divert adaptive capacity to generating disturbance rather than correcting its own mismatch, creating a self-limiting effect.
- The stochastic exponent ($b = 3/2$) is now derived from both the AR(1) stationary variance (discrete) and the Itô-Lyapunov analysis (continuous, Prop A.1S). The continuous-time analog (Ornstein-Uhlenbeck) gives the same scaling, confirming the asymptotic-scaling claim is the fluid-limit value. The 0.019 gap between the simulation $b = 1.481$ and the asymptotic $b = 3/2$ is *not pure numerical noise*: it is consistent with a derivable finite- ν correction factor (proportional to $\sqrt{(2c_{\min} - \eta_A^* c_{\max}^2)/(2c_{\min} - \eta_B^* c_{\max}^2)}$ when $\eta_A^* > \eta_B^*$) that arises because the discrete steady-state variance carries the $O(\eta^* c_{\max}^2 / c_{\min}^2) = O(c_{\max}^2 / (c_{\min}^2 \nu))$ gap from $\mathcal{O}$. In the fluid limit ($\nu \rightarrow \infty, \eta^* \rightarrow 0$ at fixed $\mathcal{T}$), the correction factor approaches 1 and the asymptotic $b = 3/2$ is recovered exactly. The two models (D and S) are unified by the common sector-condition framework with different disturbance assumptions (GA-2 vs. GA-2S).
- Asymmetric coupling ($\gamma_A \neq \gamma_B$) appears as a multiplicative prefactor γ_A / γ_B that shifts the mismatch ratio without changing the exponent. An agent with lower tempo but higher coupling effectiveness (γ) can partially compensate — but the squared dependence on tempo dominates for large tempo ratios.

Incidental audit gold (lift 2026-05-31). Cross-audit “wandering thoughts” / §14-ideation harvested from the de-novo auditors' working dirs, deduplicated across substrates and attributed by substrate + audit. Orthogonal pedagogical / framing / figure / naming material, kept separate from the certified theory-fix findings (the 526815 F215–F221 stream — the $\alpha = \mathcal{T}$ bridge dependency, the Model-S additive-noise-scale convention, undeclared proof-source dependencies, and the det-vs-stoch non-coupling-dominant limit mismatch with the intro — is certified-track and routed there, not here). **Coverage:** dedicated reflections from 193847, 829314, 849201 (Gemini), the 526815 first-pass, and the 451729/471203 Section-III batch-reflections; substrate attribution inferred from voice where not explicit.

1. Candidate Brief prose / pre-prose.

- The squared-scaling consequence, stated for a Brief: a 2:1 tempo ratio is not a 2:1 advantage but a 4:1 mismatch ratio — “the faster agent doesn't just have half the error; it inflicts four times the chaos it absorbs” (Gemini, AUDIT-WORKING-829314). The body has the numbers; this is the plain compounding gloss.
- Why $b = 2$ rather than additive, in one line: speed “simultaneously increases your defense (mismatch correction) and your offense (disturbance generation), so the two factors multiply” (Gemini, AUDIT-WORKING-849201; Gemini, AUDIT-WORKING-193847).

2. Candidate Discussion.

- **Directional-vs-noisy strategic doctrine (strongest framing here).** The $b = 2$ (Model D) vs $b = 3/2$ (Model S) split is read as an operational doctrine worth surfacing in Discussion: “don't just be fast and noisy — be fast and directional.” Randomized features / chaotic probing scale at the weaker $3/2$ power; a coordinated suite that moves the environment in a *consistent* direction scales at the full squared power, because it forces the opponent to spend its structural reserve R updating in one direction rather than being jiggled around a noise ball (Gemini, AUDIT-WORKING-829314; the segment's own “consistent directional pressure is more effective per unit tempo” Discussion line is the same point). A candidate sharpening that turns the exponent split into a prescription. (*Early-conflation texture: stated as derived doctrine; the exponents are exact-conditional, the strategic reading is interpretation.*)
- **Competitive-investment / monopoly-formation reading.** Under adversarial coupling a 10%-faster firm does not get a 10% edge — in Model D it gets $(1.1)^2 = 1.21$, a 21% structural-coherence advantage; a 2:1 tempo edge gives one-quarter the error and 4× the inflicted disturbance. Offered as a structural account of why tech monopolies form quickly: a slight tempo edge triggers the squared-feedback widening until the slower firm crosses its destabilization threshold (#der-adversarial-destabilization) and collapses (Gemini, AUDIT-WORKING-829314). A candidate Discussion instantiation of the coupling-dominant regime; the destabilization-threshold hand-off is already in the body.

3. Follow-up items.

- **Boyd's-law scope-honesty (candidate scope statement).** A break-protocol adversarial reading pressed whether the squared $3/2$ laws over-generalize Boyd's specific commander-vs-commander setting: many adversarial dynamics aren't tempo races (cybersecurity = response-time to specific threats; markets = tempo advantage up to a saturation point; evolution = fitness-landscape navigation, not a tempo race). Verdict was scope-narrowing, not refutation: the laws apply to the *coupling-dominant loop-loop class*; other classes scale differently (Claude, AUDIT-WORKING-471203, Challenge 11). The body's Regime-3 ($b \rightarrow 1$) limit and the "whether this mechanism is dominant in actual adversarial interactions is empirical" hedge already move here; this is a candidate to state the class-restriction more visibly so the OODA framing doesn't read as a universal law.
- **Placement: lift "the two effects compound" into Formal Expression.** A batch auditor suggested the intuitive "two effects compound" explanation should move into the Formal Expression as a brief derivation note, since the mechanism is non-obvious and deserves emphasis (Claude, AUDIT-WORKING-451729 batch-14). A staging-for-pedagogy nudge; the explanation is currently in the lead paragraph and Discussion.
- **The finite- ν reconciliation as a trust-compounding exemplar.** The Working-Notes treatment of the 0.019 gap (simulation $b = 1.481$ vs asymptotic $3/2$) as a *derivable* finite- ν correction factor rather than dismissed noise was independently called out across substrates as a high-water mark of the framework's mathematical hygiene — "willing to derive a correction factor rather than wave away discrepancies" (Claude, AUDIT-WORKING-471203; Gemini, AUDIT-WORKING-829314; Gemini, AUDIT-WORKING-849201). Worth preserving prominently (and a candidate for a Brief/README "how the framework treats its own discrepancies" example) when this segment is promoted.

4. Readers often ask / wonder.

- **Does the finite- ν correction mean discrete-time agents (LLMs taking turns) pay an intrinsic penalty vs continuous-time agents at equal effective tempo?** And does it follow that high-frequency small updates are always superior to low-frequency large ones? A natural reader extrapolation from the correction-factor WorkingNote (Gemini, AUDIT-WORKING-193847).
- **If both agents are fast enough to perfectly predict each other, do the γ coupling coefficients collapse?** I.e. what happens to the advantage under symmetric mutual-predictability / decoupling? (Gemini, AUDIT-WORKING-849201).

5. Candidate figures.

- **Two-factor product diagram.** The squared advantage drawn as a two-factor product: faster correction on the attacker's side $\times$ faster disturbance-generation against the target; the stochastic branch replaces one full correction factor with a square-root averaging factor, making $b = 3/2$ visually distinct from $b = 2$ (Claude, AUDIT-WORKING-526815 first-pass).

Belongs elsewhere.

- **Future-AI bounding reading (high-application reach).** The squared scaling is read as an argument that a large tempo asymmetry between a future AI and slower agents in a coupling-dominant adversarial regime is a *regime change*, not a quantitative gap: a 10:1 tempo ratio is a 100:1 mismatch ratio, almost certainly past the slower agent's reserve R and into structural collapse — so safe infrastructure must bound *both* tempo $\mathcal{T}$ and the coupling γ (how much one agent's actions can disturb others' shared environment), since bounding only one leaves the superlinear product unbounded (Gemini, AUDIT-WORKING-193847; Gemini, AUDIT-WORKING-829314). Aspirational consciousness-infrastructure / safety reach pointing at 04-eli-core/ and the protection-strategy program, not this segment's formal content. (Per project voice: stated as "future AI" with no capability-comparative claim beyond the tempo-ratio arithmetic.)

15.5 Discussion: Additional Implications & Discussion

The chapter is where AAT's signed-coupling structure first comes into formal view. Team persistence and adversarial destabilization turn out to be the same inequality viewed in opposite directions — both fall out of the sector-persistence template with effective disturbance decomposed into environmental and coupling-induced terms — and the chapter's

headline result (superlinear adversarial tempo advantage) follows from the template’s scaling without separate derivation. The recipient-side four-regime classification gives the chapter its distinctive contribution: what the emitter-side scalar $\gamma_A \mathcal{T}_A$ collapses, the recipient-side decomposition recovers, with three independent boundary conditions in B ’s existing AAT quantities and operational consequence for what kind of repair the receiving agent needs. None of the chapter’s segments carries a ## Findings section, so the implications segment is their canonical catalog home; the work is mostly synthesis: surfacing how the chapter’s machinery composes with Ch.2’s closure-defect framework and what the equilibrium hand-off to Ch.5 looks like.

Cooperative and adversarial are one signed coupling, not two theories.

The chapter intro #cooperative-adversarial-intro flags it as a chapter-level architectural commitment, and the chapter’s segments deliver on it: cooperative coupling and adversarial coupling are not separate dynamics with their own machinery. They are the *same* signed-coupling structure with γ of opposite sign. #der-team-persistence’s composite-persistence condition decomposes effective disturbance for sub-agent i as $\rho_i^{\text{eff}} = \rho_{i,\text{env}} + \sum_j \gamma_{j \rightarrow i}^{\text{adv}} \mathcal{T}_j - \sum_j \gamma_{j \rightarrow i}^{\text{coop}} \mathcal{T}_j$ — adversarial contributions enter as positive terms (effective disturbance amplification), cooperative contributions as negative terms (effective disturbance reduction). The signed-coupling form is what makes cooperative coupling *quantitatively comparable* to adversarial coupling rather than addressing them in separate vocabularies, and it is what lets the framework’s persistence template (Part I Ch.4 implications, §4) absorb both regimes as instantiations of a single Lyapunov argument.

The implications-side payoff is the duality between persistence and destabilization. #der-adversarial-destabilization makes this precise: B destabilizes iff $\rho_B^{\text{eff}} > \alpha_B R_B$, where ρ_B^{eff} includes the coupling-amplified term $\gamma_A \mathcal{T}_A$. *This is the negation of B ’s persistence condition under the same coupling model.* The two are the same inequality, read in opposite directions: persistence asks whether the inequality holds, destabilization asks whether it fails, and the coupling parameter $\gamma_A \mathcal{T}_A$ is what determines which side of the threshold the composite sits on. The framework’s posture is therefore that adversarial attacks are *not* a separate kind of phenomenon requiring separate machinery — they are the same machinery, with the attacker’s contribution entering the target’s effective disturbance via the signed-coupling structure.

The closed-form unification this enables, surfaced in Part III Ch.2 implications, deserves to be repeated as it bites here: the four prior persistence results — single-agent ($\gamma = 0$), team-persistence ($\gamma < 0$), adversarial-destabilization ($\gamma > 0$), coordination-failure ($C > \alpha$, Brooks’s Law) — are signed special cases of one inequality. Cooperative and adversarial coupling sit at opposite ends of the same γ -axis; the framework’s machinery does not pick a side, and the practitioner who understands one regime understands the other by sign change.

Universal taxonomy of inter-agent events — what the recipient sees.

#der-interaction-channel-classification is the chapter’s distinctive AAT-internal contribution: the same signal from emitter A lands on recipient B as one of *four* qualitatively different things — *informative update*, *magnitude shock*, *structural shock*, or *ambient noise* — and which one is determined by three independent boundary conditions stated entirely in B ’s existing AAT quantities. The emitter-side scalar $\gamma_A \mathcal{T}_A$ that previous segments have used as a stand-in for “how strong the coupling is” loses this distinction; the recipient-side classification recovers it. The implication for the framework’s repair machinery is concrete: the four regimes route to four different corrective paths, and “more tempo” works for some regimes but not for others.

The three boundary conditions deserve to be named even at the implications-segment grain. The first separates informative updates from magnitude shocks: whether the residual δ_τ^B is within the current model’s prediction-confidence interval or outside it. Inside the interval, the event is informative — B ’s standard gain-principle update absorbs it with full effect. Outside the interval, the event is a magnitude shock — the gain principle still applies, but the residual is large enough to suggest model-class strain rather than parameter drift. The second separates magnitude shocks from structural shocks: whether the model class $\mathcal{M}$ can absorb the event at all. When the residual reflects a regime the model class cannot represent (#result-structural-adaptation-necessity applied locally), the event triggers structural adaptation rather than parametric update. The third separates ambient noise from informative updates: whether the event’s

expected magnitude clears the agent’s signal-detection threshold; below it, the event is filtered as noise and produces no update.

The framework’s prescription is that *repair path depends on regime*, not on event magnitude or coupling strength alone. Informative updates respond to tempo investment (more $\mathcal{T}_B$ closes the residual faster). Magnitude shocks respond to gain investment (lower $U_{o,B}$ — better observation quality — closes more of each shock per cycle). Structural shocks respond to model-class expansion (#result-structural-adaptation-necessity’s structural change) — no amount of tempo or gain suffices. Ambient noise calls for *infrastructure-level filtering* below the agent proper (#der-interaction-channel-classification’s Regime-III repair): the events produce no usable update, but processing them consumes tempo and their variance drains adaptive reserve, so the repair is to keep them from reaching the core loop at all — not to invest agent-level detection in them. The implication for adversarial robustness: an attacker who knows the target’s regime boundaries can choose events that *miss* the target’s repair vocabulary — structural shocks against a tempo-investment defender, magnitude shocks against a model-class-expansion defender. The framework’s design prescription is that defenders should *cover all four repair paths*, not optimize for one.

The composite-level payoff is the regime-typed effective disturbance decomposition. The chapter’s ρ_B^{eff} decomposes per regime, matching the displayed decomposition in #der-interaction-channel-classification: magnitude shocks contribute peak-load disturbance; structural shocks contribute the structural mismatch floor and set structural-adaptation cadence; ambient noise contributes *variance* — it enters Model S’s $\sigma_{w,B}^2$ below the observability floor, yielding no usable update while slowly draining adaptive reserve (the drain-without-information is exactly why a Regime-III flood is an attack rather than an inconvenience); and informative updates enter with *negative* sign — they reduce effective disturbance, generalizing #der-team-persistence’s cooperative term. The aggregate effective-disturbance number from #der-team-persistence is the signed *sum* across regimes; the per-regime decomposition is what tells the practitioner *where the disturbance is coming from*. For composite-agent design under uncertainty about the disturbance regime, the framework’s prescription is to instrument the four regime detectors independently — knowing ρ_B^{eff} is small is less useful than knowing which regime is contributing what.

Adversarial tempo advantage as superlinear scaling — and the disturbance-model split.

#result-adversarial-tempo-advantage carries the chapter’s headline result, and the implication worth surfacing is *why* the scaling is superlinear. The template’s Model D conclusion gives $\|\delta\|_{ss} = \rho^{\text{eff}}/\mathcal{T}$ in the linear case; substituting the coupling-amplified disturbance gives the steady-state mismatch ratio scaling as $(\mathcal{T}_A/\mathcal{T}_B)^2$ — superlinear by exponent $b = 2$ under deterministic drift coupling. Under Model S (stochastic) the template’s $1/\sqrt{\alpha}$ scaling produces $b = 3/2$. The two exponents are not arbitrary parameters: they fall out of Part I Ch.4’s persistence-template structure, $1/\alpha$ versus $1/\sqrt{\alpha}$, applied to the coupling model. The framework’s posture is that *the same template that explains why adaptive systems persist also predicts the precise rate at which they lose to a faster attacker* — the dual is structural.

The exponent splits further in Ch.5 — #result-adversarial-exponent-regimes gives three regime-typed values ($b = 2, 3/2, \sim 1$), with the coupling-non-dominant case dropping to $b \rightarrow 1$. The non-superlinear regime is operationally important: when observation noise is high enough that the attacker’s tempo advantage doesn’t penetrate the defender’s observation channel cleanly, the superlinear scaling collapses to roughly linear, and the framework predicts that *cheap noise injection* into the defender’s observation channel actually *helps* the defender against a high-tempo attacker by reducing the exponent. This is the structural reason the empirical pattern that highly-noisy environments produce more cooperative-game-like dynamics than expected — high observation noise gates the adversarial tempo advantage to near-linear scaling, where the attacker cannot exploit superlinear scaling to drive the defender past the destabilization threshold. The full regime taxonomy lives in Ch.5; the chapter here establishes that superlinear *is* the dominant regime under coupling-dominant noise structures.

The operational anchor worth carrying is Boyd’s OODA-loop aphorism, which the framework now gives precise mathematical content. *Inside the opponent’s loop* means $\mathcal{T}_A > \mathcal{T}_B/k$ for some regime-dependent k , and what the inside-the-loop position buys the attacker is superlinear destabilization rate rather than mere tempo advantage. The chapter’s intro framing — “superlinear tempo advantage as regime change rather than quantitative gap” — is the implications-flavored

point: at the destabilization threshold, the target does not gradually degrade. It loses bounded behavior, the way a balance held just barely beneath a tipping point is qualitatively different from one well above it. The threshold is the same one Part I Ch.4's central inequality names; the chapter's contribution is identifying the coupling-induced disturbance term that pushes the target past it.

Three adversarial obstructions to contraction — the equilibrium hand-off.

#result-contraction-template (Appendix A, depended on by Ch.4 via the persistence template) names the chapter's most consequential structural recognition: contraction-metric machinery *cannot* handle strategic / adversarial regimes for three convergent reasons that are individually classical results from different mathematical literatures. *Slotine 2003 compositional applicability fails*: saddle-point equilibria break the attracting-fixed-point assumption that contraction analysis requires. *Passivity universality fails*: adversarial inputs drive any storage function, so passivity-based stability does not give bounds. *Daskalakis et al. 2018 last-iterate non-convergence*: for generic non-zero-sum games, learning dynamics fail to converge to equilibrium in the last-iterate sense — only in time-average sense.

The implications-side payoff is structural rather than parametric: three *unrelated* mathematical obstructions converge on the same conclusion. Contraction analysis works for cooperative / coordination regimes (where there is an attracting fixed point) and fails for strategic / adversarial regimes (where equilibrium machinery is required). The convergence of three obstructions makes the limit *structural* rather than framework-specific — it is not that AAT's contraction machinery is too weak; it is that the regime is outside the scope condition contraction analysis depends on across multiple formal traditions. The framework's posture is to recognize the limit precisely and hand off cleanly to equilibrium-theoretic methods at the boundary.

The hand-off lands in Ch.5: #deriv-strategic-composition is where equilibrium machinery (potential games via Monderer-Shapley 1996, monotone games via Rosen 1965) takes over from contraction analysis for the symmetric partially-opposing case. The chapter's contribution is naming the limit; Ch.5's contribution is supplying the replacement machinery. The framework's posture across the boundary is that *scope-honesty is load-bearing architecture*: knowing where each tool applies cleanly is more useful than trying to extend one tool across all regimes. The three convergent obstructions are why this matters operationally — an AAT-style sector-condition argument applied to a strategic regime would not just produce wrong predictions, it would produce *systematically* wrong predictions in each of the three directions the obstructions name.

The chapter also surfaces a quieter consequence worth carrying: cooperative coupling that is *too strong* can produce equilibrium-stability failures even without adversarial components. When sub-agents' cooperative coupling drives them toward a saddle-point equilibrium (rather than a stable attractor), the framework's prediction is the same as the strategic case — contraction analysis fails, and equilibrium machinery is required. This is empirically the situation in tightly-coupled multi-agent ML systems where reward-shaping creates implicit competition among cooperators; the framework predicts the failure mode is structurally identical to the adversarial case, with the unique broadly-available remedy being explicit equilibrium-theoretic analysis rather than tempo or gain investment.

Closing the chapter and bridging to Strategic Composition.

The chapter delivers the signed-coupling structure that the rest of Part III operates on. Cooperative coupling reduces effective disturbance; adversarial coupling amplifies it; the same template handles both via the sign of γ . The four-regime taxonomy says what the recipient actually sees, and the framework's repair vocabulary diversifies accordingly. The super-linear adversarial tempo advantage falls out of the template's $1/\alpha$ versus $1/\sqrt{\alpha}$ scaling; OODA's "inside the opponent's loop" acquires mathematical content. The three convergent obstructions to contraction analysis name precisely where the chapter's machinery hands off to equilibrium-theoretic tools.

What the chapter does *not* deliver is the equilibrium machinery itself, the agent-opacity formalism, the recipient-side classification's emitter-side dual, or the per-dimension persistence treatment that adversarial targeting depends on. Those

all live in Ch.5, where the framework’s machinery operates at full Part III scope: strategic composition with partially-opposing objectives, agent opacity as the formal dual of observation quality, the 16-cell emitter-recipient adversarial-targeting decomposition, the regime-typed exponent decomposition ($b = 2, 3/2, \sim 1$), and the matrix-Loewner persistence formulation that adversarial concentration on weak dimensions requires. The chapter-end here is the last single-layer-of-coupling synthesis before Ch.5’s strategic-equilibrium machinery takes over.

The bridge into Ch.5 runs through three threads. First, the contraction-template obstructions hand off to equilibrium machinery (potential / monotone games). Second, the recipient-side classification will compose with the emitter-side opacity classification in Ch.5 into a 16-cell joint composition with closed-form arg-max for adversarial-targeting decisions — the chapter here supplies one half, Ch.5 supplies the other. Third, the signed coupling γ structure becomes the *trust-meta-model signal* via Ch.3’s #hyp-communication-gain: observed cooperative coupling drives $U_{\text{align},ji} \rightarrow 0$ (high trust calibration); observed adversarial coupling drives $U_{\text{align},ji}$ large (low trust calibration). The implications-segment series continues to weave back across earlier chapters, with each chapter’s machinery operating on what previous chapters have built.

Working Notes.

- This segment is a chapter-end discussion grouping findings from #der-team-persistence, #der-adversarial-destabilization, #der-interaction-channel-classification, #result-adversarial-tempo-advantage, #cooperative-adversarial-intro, #result-sector-persistence-template (Appendix A, the template these segments instantiate), #result-contraction-template (Appendix A, the contraction-obstructions surfacing), and #deriv-strategic-composition (Ch.5, the forward-pointer for equilibrium machinery).
- **Cross-reference vs canonical-home policy.** None of this chapter’s source segments carry ## Findings sections. The implications segment is the canonical catalog home for: #21 (Universal Taxonomy of Inter-Agent Events — #der-interaction-channel-classification content), the cooperative-vs-adversarial duality (the signed-coupling structure as one architectural commitment), the headline adversarial tempo advantage result (#44 partial — the $b = 2$ and $b = 3/2$ exponents from this chapter; the regime-decomposition refinement lives in Ch.5), and #24 (Three Adversarial Obstructions to Contraction Analysis — #result-contraction-template content surfaced at this chapter’s hand-off boundary).
- **Forward-pointers into Ch.5 are heavy.** The chapter sits structurally between cooperative-coupling machinery (here) and strategic-composition machinery (Ch.5). Several findings are partially treated here with completions in Ch.5: regime-typed exponent decomposition, observation-noise gating of advantage, agent-opacity formalism as dual to observation quality, 16-cell emitter-recipient adversarial-targeting decomposition, and the matrix-Loewner persistence formulation that adversarial concentration depends on. Reading this implications segment together with Ch.5’s gives the full coupling picture; reading either alone gives the half-picture.
- **Identifiability-floor pattern coverage update.** No new instances surface in this chapter — the four formal instances (Instances 1–4) were all in view by Part II Ch.5’s implications, and Ch.5 here will give Instance 3 (composite contraction certification from component marginals via Liberzon common-Lyapunov-nonexistence) its operational form in the strategic-composition setting. The contraction-obstructions in #result-contraction-template are *adjacent* to but not the same as the identifiability-floor pattern — they are obstructions to a specific *analysis method* rather than no-go theorems about *identification*. The distinction is worth carrying: identifiability-floor instances forbid identifying a quantity from data; contraction-obstructions name where a method fails. Both are scope-honesty-load-bearing for AAT, but they sit in different meta-architectural slots.
- **Connection to NeurIPS submissions.** The superlinear tempo advantage result has not been formalized to theorem grade in current submissions; the signed-coupling structure underlies Paper 2’s multi-agent extension (which is a candidate for a future revision rather than current scope). The chapter’s content stays in the AAT-internal register.

Incidental audit gold (lift 2026-05-31). Of the de-novo audit dirs, only the deeply-mathematical certified-track auditor (526815) reached a dedicated reflection on this chapter-end implications segment; the gold-stream Gemini/Claude dirs did not file a separate reflection on it. **No framing/pedagogy/figure/naming gold surfaced** — and the certified-track signal that *did* surface is a single coherent theme worth flagging for the findings channel rather than the gold channel:

Off-ramp (NOT gold — flagged for the certified track).

- **Synthesis-drift: the implications segment firms up caveats its source segments stated carefully (526815 F222–F232).** As a chapter-end catalog that imports conditional and discussion-grade pieces from #der-team-persistence / #der-adversarial-destabilization / #der-interaction-channel-classification / #result-adversarial-tempo-advantage plus forward Ch.5 material, this segment was read as repeatedly converting source caveats into stronger operational prescriptions. The specific drifts flagged, each a candidate to reconcile against the weaker source status: (i) *(executed 2026-07-03, audit 731548 B-4 — the body’s regime-typed decomposition and repair-path paragraphs now match the source’s displayed formula, including the Regime-III variance drain and the negative Regime-I sign; F224/F225 closed)*; (ii) **magnitude-shock repair stated as gain/ U_o investment** here vs the classification segment’s capacity/sector-radius axis (F223); (iii) **“contraction-metric machinery cannot handle strategic/adversarial regimes” as a class claim** is broader than the source’s scoped attracting-fixed-point obstruction (F229), and the companion “adversarial inputs drive any storage function” passivity claim is stated too sweepingly without the specific input/output passivity assumptions (F230); (iv) **“inside the opponent’s loop means $\mathcal{T}_A > \mathcal{T}_B/k$ ”** is not the threshold the derivations deliver — destabilization depends on $\gamma_A \mathcal{T}_A$ vs B ’s reserve, the tempo-ratio result needs the coupling-dominant symmetric limit (F228); (v) **“cheap noise injection into the defender’s observation channel helps the defender”** is one-sided — higher U_o also degrades the defender’s real observations, update gain, and ordinary persistence (F226), and the attached “noisy environments produce more cooperative-game-like dynamics” is an unsupported cross-domain claim that conflates adversarial-channel noise with shared-task noise (F227); (vi) the **trust-calibration bridge** ($U_{\text{align}} \rightarrow 0$ / large from observed coupling sign) maps coupling sign directly to alignment certainty, where it should be *evidence* updating trust under a model (F232). This is the recurring “implications segments must preserve the weakest status of each imported claim” discipline (the segment’s own Working Notes already cite the cross-reference-vs-canonical-home policy); routed to the certified-findings track for adjudication, recorded here because it is a substantive epistemic-status pattern across the whole segment, not framing gold.

Chapter 16

Strategic Composition and Channel Effects

[Gap] Discussion

16.1 *Derivation*: Strategic Composition via Equilibrium Convergence

When two or more AAT agents interact through a shared environment with **partially-opposing objectives** $\{O_t^{(i)}\}$, the composition-level question is not “does the trajectory contract to zero closure-defect?” but “does the coupled best-response dynamics admit an equilibrium and converge to it?” Contraction to shared truth is a $U_O = 1$ special case; strategic composition is the $U_O < 1$ companion regime in which the correct primitive is **fixed-point existence and stability**, not Lyapunov contraction on a shared state. The framework names three composition-level questions distinctively here: **(SC-1)** does an equilibrium joint policy exist where no sub-agent has a unilateral mismatch reduction available; **(SC-2)** do coupled best-response dynamics initialized near it stay there; **(SC-3)** do they converge from arbitrary initialization to the equilibrium set? Fixed-point, local-stability, and reachability questions respectively — none a Lyapunov contraction question on a shared state variable.

Two sub-scopes carry progressively weaker conditions. Under potential-game (Monderer-Shapley 1996) or monotone-game (Rosen 1965) conditions — *sub-scope* α' — the sector-persistence template transfers to the gradient of the joint potential (resp. to a weighted-norm variational inequality on the joint Jacobian’s symmetric part), and AAT’s persistence machinery recovers at the equilibrium layer with state variable $\xi = \pi - \pi^*$ (deviation from Nash) and sector constant α_{joint} living at the *joint potential’s curvature* rather than at any individual sub-agent’s α . Outside α' — *sub-scope* β' — only set-convergence to coarse correlated equilibria is available (Hart-Mas-Colell 2000, rate $O(1/\sqrt{T})$), and the macro-state of a strategic composite is a *distribution* on joint strategy space rather than a state-space point. The right structural consequence is named on the **dynamic-regime axis** rather than the architectural-class axis (see #disc-dynamic-regime-axis): the strategic-composition move transitions the composite from contraction-regime (R0, under aligned objectives — scope routes C-i/ii/iii) to equilibrium-regime (R1, under partially-opposing objectives with potential/monotone structure — sub-scope α') or to cyclic-distributional-regime (R2, under partially-opposing objectives without potential/monotone structure — sub-scope β'). The composite’s *macro-state type* changes correspondingly — state-variable in R0 / fixed-point object in R1 / distributional object in R2 — but the *architectural class* of the composite is preserved under goal-blind routing and distinct sub-agent substrates (it is Class 1 per the formal criterion when the sub-agents are Class 1; the Class 2 / Class 3 changes per #hyp-directed-separation-under-composition Case 2 require routing-structure goal-dependence or shared-substrate G^c -allocation, not strategic composition itself). This segment establishes the framing, derives the α' -transfer (R0 $\rightarrow$ R1 Lyapunov-machinery transfer per #disc-dynamic-regime-axis

§"Lyapunov-machinery transfer $R0 \rightarrow R1$ "), documents the β' scope limits honestly, and relates to #der-adversarial-destabilization (asymmetric adversarial) and #deriv-critical-mass-composition (aligned composition with shared target) as siblings on the composition axis.

Formal Expression.

The framing move.

For N purposeful sub-agents with partially-opposing $\{O_t^{(i)}\}$ running coupled AAT loops through a shared environment, the composition-level question is:

Formulation (strategic-composition-framing).

(SC-1) Existence of equilibrium. Does there exist a joint state $(X_{c,1}^*, \dots, X_{c,N}^*)$ — equivalently a joint policy profile $(\pi_1^*, \dots, \pi_N^*)$ — such that no sub-agent has a unilateral mismatch reduction available?

(SC-2) Stability of equilibrium. Do coupled best-response dynamics, initialised near $(X_{c,i}^*)$, remain there or return to it?

(SC-3) Convergence from interior. Do coupled dynamics, initialised away from any equilibrium, converge to the equilibrium set?

(SC-1)–(SC-3) are a fixed-point-existence question, a local-stability question, and a reachability question respectively. None is a Lyapunov contraction question on a shared state variable. The contraction framing in #form-composition-closure / #deriv-critical-mass-composition is recovered as the $U_O = 1$ special case (unique Nash equilibrium at the shared-objective optimum; best-response dynamics collapse to single-objective optimization).

Sub-scope α' : potential and monotone games (A2'-analog derived).

A strategic interaction $(\{A_i\}, \{O_t^{(i)}\})$ is a **potential game** (Monderer & Shapley 1996, *Games and Economic Behavior* 14) when there exists a scalar potential Φ such that each sub-agent's unilateral improvement matches the potential's unilateral improvement:

Derived (A2'-analog-potential, from Monderer-Shapley 1996).

$$O_t^{(i)}(\pi'_i, \pi_{-i}) - O_t^{(i)}(\pi_i, \pi_{-i}) = \Phi(\pi'_i, \pi_{-i}) - \Phi(\pi_i, \pi_{-i}) \quad \forall i. \quad (16.1)$$

Under this condition plus each sub-agent's B1 directional fidelity (4.3) to $\nabla_{\pi_i} O_t^{(i)}$, the joint best-response dynamics satisfy

$$\frac{d\Phi(\pi)}{dt} = \sum_i \langle \nabla_{\pi_i} \Phi, \dot{\pi}_i \rangle \geq \alpha_{\text{joint}} \|\nabla \Phi\|^2 \quad \text{for some } \alpha_{\text{joint}} > 0 \quad (16.2)$$

whenever the joint configuration is not at a stationary point of Φ . This is (T2) transcribed to state variable $\xi =$ gradient-of-potential, correction function $F =$ joint best-response velocity field; the quadratic Lyapunov structure is the same. **The sector-persistence template transfers** with $\xi = \pi - \pi^*$ (deviation from Nash), $\alpha = \alpha_{\text{joint}}$, $R =$ basin-of-attraction radius, and ρ_ξ the exogenous disturbance rate on strategy space. Equilibrium existence follows from potential-function compactness on compact strategy spaces; equilibrium stability follows from Φ 's role as a joint Lyapunov function.

Weaker than potential: a game is **diagonally strictly concave** (Rosen 1965, *Econometrica* 33) when the Jacobian of the joint gradient field is negative-definite on the joint strategy space. Under this condition, there exists a unique Nash

equilibrium and the joint gradient dynamics converge to it exponentially. The convergence rate is bounded below by the smallest eigenvalue of the symmetric part of the joint Jacobian, playing the role of α_{joint} . No scalar potential need exist; a *weighted-norm* Lyapunov argument on the joint Jacobian’s symmetric part substitutes.

Derived (A2'-analog-monotone, from Rosen 1965).

Sub-scope α' comprises: potential games (Monderer-Shapley), monotone games (Rosen), strongly-monotone games, and exponential-family dual-averaging under concave objectives. For these classes, the sector-persistence template extends to equilibrium convergence with composite sector constant α_{joint} inheriting from the joint-gradient-field structure.

Sub-scope β' : non-potential non-monotone games.

Every strategic interaction with continuous strategy spaces and regular payoffs can be reformulated as a **variational inequality**: find $\pi^* \in \mathcal{K}$ such that $\langle F(\pi^*), \pi - \pi^* \rangle \geq 0$ for all $\pi \in \mathcal{K}$, where F is the joint pseudo-gradient field. When $\mathcal{K}$ is compact-convex and F is continuous, a solution exists (Hartman-Stampacchia theorem). **Pure-strategy Nash equilibrium existence is therefore guaranteed** for continuous-strategy games with compact convex strategy sets and continuous payoffs. But the VI framework gives *existence* only, not *convergence of any specific dynamic to the solution*; solutions may be non-unique.

Derived (equilibrium-existence-via-VI, from Facchinei-Pang 2003).

Under no-regret learning (e.g., Hedge / multiplicative weights, Freund-Schapire 1997), the empirical joint distribution converges to the set of **coarse correlated equilibria** (CCE) at rate $O(1/\sqrt{T})$. This requires no structure on the game beyond each sub-agent computing its own regret.

Derived (regret-minimization-convergence-to-CCE, from Hart-Mas-Colell 2000).

Under β' , the macro-state of a strategic composite is a *distribution* on the joint strategy space — the empirical joint play whose support is the CCE — rather than a state-space point. This is the structural shape of “convergence” in the β' regime: distributional convergence, not pointwise. Pure-strategy Nash may or may not exist (cyclic games — rock-paper-scissors, matching pennies — lack pure Nash but retain mixed Nash via Nash 1950 and CCE convergence via Hart-Mas-Colell 2000); the β' machinery’s guarantees are at the distributional layer regardless.

Sub-scope β' scope-honesty: AAT can predict that long-run joint play lies in the CCE support; it cannot predict short-run trajectory, per-sub-agent mismatch convergence, or selection among multiple equilibria. The sector-persistence template does *not* apply in β' . This is a genuine scope limit shared with game theory as a whole, not a defect of AAT.

Cross-layer linearization fingerprint. The joint Jacobian of F at saddle-only Nash equilibria in β' has imaginary-axis (semisimple) spectrum — the linearization fingerprint of the R0-loss rung in the single-agent certificate-strength ladder of #result-certificate-existence. FTRL $\times$ graphical-constant-sum-with-fully-mixed Nash (Cheung-Piliouras-Tao 2021 Theorem 19: lossless DGS + Poincaré recurrence on bounded level sets) is the worked composite instance of R0-loss; the connection is between layers (single-agent linearized error space $\mathbb{R}^n$ vs composite joint-strategy space $\mathcal{X}^c$), not a collision of R-letter ladders.

Zero-sum scalar worked example.

Two agents A, B with scalar actions $a_i \in [-1, 1]$ and state $s_{t+1} = s_t + a_A - a_B + w_t$, $w_t \sim \mathcal{N}(0, \sigma^2)$. Objectives $O_t^{(A)}(s) = s$ (maximize s), $O_t^{(B)}(s) = -s$ (minimize s); zero-sum at the state-dependent payoff level.

Derived (zero-sum-scalar-instantiation).

Action-coefficient analysis. The agents’ state-preferences are opposite, but the action-coefficients in s_{t+1} are also opposite: $\partial s / \partial a_A = +1$ and $\partial s / \partial a_B = -1$. Composing with the objectives’ state-dependence:

$$\frac{\partial O^{(A)}}{\partial a_A} = \frac{\partial O^{(A)}}{\partial s} \cdot \frac{\partial s}{\partial a_A} = (+1)(+1) = +1 \quad (16.3)$$

$$\frac{\partial O^{(B)}}{\partial a_B} = \frac{\partial O^{(B)}}{\partial s} \cdot \frac{\partial s}{\partial a_B} = (-1)(-1) = +1 \quad (16.4)$$

Both agents marginal-prefer *increasing* their own action — the opposing state-preferences and opposing action-coefficients compose to aligned action-preferences.

Potential function. For Monderer-Shapley, $\partial\Phi/\partial a_i = \partial O^{(i)}/\partial a_i$ for each i . With both partial derivatives equal to +1, the potential is $\Phi(a_A, a_B) = a_A + a_B$ (modulo additive constant).

Nash equilibrium. Each agent's best response on $[-1, 1]$ is to push its own action to +1; the unique Nash equilibrium is $(a_A^*, a_B^*) = (1, 1)$. At equilibrium, $\Delta s = a_A^* - a_B^* = 0$: the substantive zero-sum property is that equal opposing action-coefficients cancel under the agents' aligned action-preferences, so joint action contributes no net displacement to s (state drift comes only from w_t). This is qualitatively different from the naive picture in which the agents push the state in opposite directions.

Sector-persistence template — interior-NE instantiation via Cournot-style payoffs. The unregularized example has a *linear* potential Φ and a *corner* equilibrium at $(1, 1)$; the unconstrained gradient field $\nabla\Phi = (1, 1)$ is constant in $\xi = \pi - \pi^*$, so there is no interior linear restoring force around the equilibrium and the template's (T2) sector lower bound $\xi^T F(\xi) \geq \alpha \|\xi\|^2$ does not hold without modification — the saturation, not a sector contraction, is what enforces equilibrium. The aligned-action-preferences finding is the conceptual payload of the unregularized form; the sector-template transfer requires a payoff structure whose curvature is genuinely quadratic, not merely projected onto a corner.

A Cournot-style duopoly (Cournot 1838; Monderer-Shapley 1996 §3 example) supplies that structure naturally. Two firms produce quantities $q_i \in [0, Q_{\max}]$; market price is $P(q_A, q_B) = a_0 - b(q_A + q_B)$; per-firm cost is linear at rate κ ; per-firm profit is

$$O^{(i)}(q_A, q_B) = q_i (a_0 - b(q_A + q_B) - \kappa). \quad (16.5)$$

This is partially-opposing in the policy-conflict sense — each firm's marginal payoff is decreased by the other's production through the shared price — without being strictly zero-sum. It is a potential game with quadratic potential.

Derived (Cournot-as-potential, from Monderer-Shapley 1996 §3).

$$\Phi(q_A, q_B) = (a_0 - \kappa)(q_A + q_B) - b(q_A^2 + q_B^2) - b q_A q_B, \quad (16.6)$$

verified by $\partial\Phi/\partial q_i = (a_0 - \kappa) - 2bq_i - bq_{-i} = \partial O^{(i)}/\partial q_i$. The symmetric interior Nash equilibrium is $q^* = (a_0 - \kappa)/(3b)$, located strictly inside the action box whenever $0 < q^* < Q_{\max}$.

Under continuous-time best-response $\dot{q}_i = \partial O^{(i)}/\partial q_i$, writing $\xi = q - q^* \mathbf{1}$, the joint dynamics are $\dot{\xi} = -bM\xi$ with $M = \begin{pmatrix} 2 & 1 \\ 1 & 2 \end{pmatrix}$. Setting $F(\xi) = bM\xi$ (correction-strength sign convention; $\dot{\xi} = -F(\xi)$):

- (T1) $F(0) = 0$. ✓
- (T2) $\xi^T F(\xi) = b(2\xi_A^2 + 2\xi_B^2 + 2\xi_A \xi_B) \geq b\|\xi\|^2$, since M 's eigenvalues are $\{1, 3\}$. So $\alpha_{\text{joint}} = b$ — the smallest eigenvalue of the symmetric part of the joint Jacobian, exactly the quantity Rosen 1965 identifies as the convergence rate. ✓

- (T2-upper) $\|F(\xi)\| \leq 3b\|\xi\|$ from M 's largest eigenvalue.
- (T3) $R = \min(q^*, Q_{\max} - q^*)$ — the radius of the largest origin-centered ball in ξ -space contained in the feasible region; the closer of the two active constraints (zero-production and capacity) determines the basin.

The template's preconditions hold with α_{joint} inherited from the demand-side curvature parameter b rather than from any individual sub-agent's α . This is the substantive transfer: the composite-level sector constant lives at the *joint potential's curvature*, with economic interpretation (market-saturation slope) rather than ad-hoc parametrization. Other linear-quadratic strategic-substitutes games (LQR with coupled state, network-coordination with quadratic disagreement cost, public-goods with quadratic externalities) produce structurally equivalent template instantiations, with α_{joint} tracking the symmetric part of their respective joint Jacobians.

The unregularized scalar zero-sum and the Cournot instantiation are complementary: the former illustrates the framing's surprise — opposing state-preferences combined with opposing action-coefficients yield aligned action-preferences — while the latter exhibits the template transfer with genuine quadratic curvature and interior NE.

New scope route (C-iv) in #scope-composite-agent.

Strategic composition with partially-opposing $\{O_t^{(i)}\}$ admits a joint equilibrium structure $\mathcal{E}$ (Nash, correlated, or coarse correlated) such that coupled best-response dynamics converge to the support of $\mathcal{E}$. The composite exists as an AAT agent with macro-state defined relative to $\mathcal{E}$ rather than relative to a shared target state.

Proposed Scope (composition-scope-condition, route C-iv).

(C-iv) is qualitatively distinct from (C-i)–(C-iii): it does *not* require shared objectives, hierarchical derivation, or mutual benefit. It requires only structural convergence of the strategic interaction in the game-theoretic sense. Composites satisfying (C-iv) are **strategic composites**, distinguished from alignment composites (C-i, C-ii) and mutual-benefit composites (C-iii).

What Is Derived vs. What Is Chosen.

Table 16.1

<i>Property</i>	<i>Source</i>	<i>Strength</i>
Framing move (contraction $\rightarrow$ equilibrium convergence)	Standard game-theoretic positioning	Formulation choice
(SC-1) / (SC-2) / (SC-3) three-question decomposition	Parallel to fixed-point / stability / reachability in dynamical systems	Formulation choice
A2'-analog under potential-game condition	Monderer-Shapley 1996 transcribed into AAT notation with B1 directional fidelity	Derived (exact under potential-game + B1)
A2'-analog under monotone-game condition	Rosen 1965 transcribed with weighted-norm Lyapunov	Derived (exact under diagonal strict concavity)
Equilibrium existence via VI	Facchinei-Pang 2003 (Hartman-Stampacchia)	Derived (external theorem, applied to strategic composition setting)
Regret-minimization CCE convergence	Hart-Mas-Colell 2000	Derived (external theorem applied)
Sub-scope α' / β' partition	Parallel to #deriv-sector-condition α/β	Formulation choice
Zero-sum scalar instantiation (corner-NE conceptual lesson)	Direct substitution into potential-game framework	Exact (within stated setup)
Cournot-style sector-template instantiation (interior-NE quadratic)	Monderer-Shapley 1996 §3 + sector-persistence-template; demand-side curvature b supplies α_{joint}	Exact (within stated setup)
(C-iv) scope route	Extension to #scope-composite-agent disjunction	Formulation choice (scope extension)

<i>Property</i>	<i>Source</i>	<i>Strength</i>
Effects-spiral eigenvalue condition	Joint-Jacobian $\text{Re}(\lambda_{\max}) > 0$ at candidate equilibria — formalizes #der-adversarial-destabilization's discussion-grade effects spiral	Sketch (specific AAT instantiations open)
Strategic composition transitions composite from contraction-regime (R0) to equilibrium-regime (R1, under α') or cyclic-distributional-regime (R2, under β') per #disc-dynamic-regime-axis	Game-structure (potential / monotone / cyclic) under partially-opposing objectives, <i>not</i> architectural-class change	Derived (regime-axis, surfaced 2026-05-21; supersedes the earlier "Class 1 sub-agents $\rightarrow$ Class 2 composite" architectural-class framing)
Mechanism-design impossibility landed in #disc-implementation-impossibility (designer-side sister meta-pattern to identifiability-floor)	Gibbard-Satterthwaite 1973-75, Arrow 1951, Myerson-Satterthwaite 1983	Charter instances at #deriv-strategy-proofness-impossibility / #deriv-bilateral-trade-impossibility / #deriv-social-welfare-aggregation-impossibility; sub-scope α' named as adjacent-but-not-identical to the preference-domain-restriction escape
Potential function Φ as additive-coordinate-forcing instance	Φ 's additivity is definitional consequence of being a potential game, not forced by AAT-internal axiom; adjacent family member	Discussion-grade
Bridge from strategic composition to #form-composition-closure $\mathcal{E}^*$	Macro-description is an equilibrium statistic, not an equilibrium state	Open
General equilibrium-selection under multiple Nash	Existence theorems do not pin down which equilibrium; selection (risk-dominance, payoff-dominance, Pareto) is partial	Open
Mean-field-game limit ($N \rightarrow \infty$)	Lasry-Lions 2007; Huang-Malhamé-Caines 2006; requires population scope condition	Open (pending Part III population gaps)

Epistemic Status.

Conditional. Max attainable: *exact* under potential-game + B1 (sub-scope α'); *derived* under monotone-game + diagonal strict concavity; *discussion-grade* for the framing and sub-scope β' set-convergence-only claim.

The A2'-analog results under Monderer-Shapley 1996 and Rosen 1965 are transcriptions of established game-theoretic theorems into AAT notation. AAT's contribution is not the mathematics but the recognition that the sector-persistence template transfers cleanly when these conditions hold, and that the composite-level sector constant α_{joint} lives at the *joint potential gradient* or *joint Jacobian's symmetric part*, not at any individual sub-agent's α . The framing move (contraction $\rightarrow$ equilibrium convergence) is a structural positioning move.

Sub-scope β' gives AAT substantially weaker predictive power than sub-scope α' : *set-convergence to CCE only*, not trajectory convergence or per-agent mismatch convergence. This is an honest scope limit, not a defect — it mirrors game theory's own scope at the no-potential-no-monotonicity regime. AAT does not claim to predict equilibrium selection under multiple Nash, short-run dynamics in cyclic games (rock-paper-scissors), or convergence rates better than $O(1/\sqrt{T})$ in β' .

What this segment does not establish:

- General convergence proof for non-potential games under AAT-native update rules (imports Monderer-Shapley + Rosen as scope statements, doesn't re-prove them).

- Explicit joint-Jacobian analysis for specific AAT agent architectures (e.g., two Beta-Bernoulli agents on a shared DAG). Effects-spiral formalization is sketched, not derived per-instance.
- Bridge from strategic composition to #form-composition-closure's closure-defect ϵ^* — the macro-description of a strategic composite is a distributional equilibrium statistic rather than a state, so the closure-defect formulation needs re-specification. Open.
- Full mechanism-design derivation (VCG, Bayesian-Nash) for specific AAT-relevant settings.

Honest Limits.

- **No pure-strategy Nash equilibrium** in cyclic games (rock-paper-scissors; matching pennies). Mixed-Nash equilibrium exists universally for finite games (Nash 1950) but is a saddle point of best-response dynamics rather than a basin attractor; fictitious play orbits the mixed-Nash without converging pointwise to it. No-regret dynamics still drive the *empirical joint distribution* to the CCE set at rate $O(1/\sqrt{T})$ (Hart-Mas-Colell 2000), so β' machinery applies; the macro-state of a strategic composite in this regime is a distribution over the strategy space, not a state-space point.

Scope honesty — strategic-composition-failures.

- **Multiple equilibria with ambiguous selection** (coordination games; drive-left vs drive-right). Convergence to *some* equilibrium; AAT does not predict which.
- **Slow mixing / finite-sample failure** under regret-minimization: $O(1/\sqrt{T})$ cumulative regret implies finite-horizon play may be far from CCE support.
- **Non-compact strategy spaces** (unbounded prices, resource allocations). No equilibrium without compactification; compactification may be artificial.
- **Bayesian games with strategic + epistemic uncertainty compounding.** Existence and uniqueness become pathological under certain information structures.
- **Mean-field games** ($N \rightarrow \infty$) require population-scope machinery not covered here.

Discussion.

Sibling to #der-adversarial-destabilization. #der-adversarial-destabilization handles the *asymmetric* adversarial case: one agent is target; attacker's tempo is exogenous parameter; sector-persistence template applies to target's mismatch with coupling-amplified disturbance. This segment handles the *symmetric* strategic case: both agents running full AAT loops; state variable is joint deviation from equilibrium; correction function is joint best-response field. The two are siblings. #der-adversarial-destabilization's Working-Notes-flagged "coupled Lyapunov analysis is the open problem" has its formal home here — but the analysis is *not* a Lyapunov problem; it is a fixed-point problem, and the correct primitive is equilibrium convergence rather than Lyapunov descent. #der-team-persistence is recovered as the cooperative limit (all γ^{coop} dominate; joint dynamics reduce to parallel single-agent sector-persistence with reduced effective disturbance); #deriv-critical-mass-composition sits in the middle with signed γ on a matched-symmetric-Tier-1 dyad under shared target.

The effects spiral's formal home. #der-adversarial-destabilization's effects spiral — B 's degrading model causes B 's actions to destabilize A further — is a **joint-Jacobian eigenvalue condition** in the strategic-composition framing: the spiral exists iff $\max_{\text{candidate equilibria}} \pi^* \text{Re}(\lambda_{\max}(\nabla F(\pi^*))) > 0$, where F is the joint best-response field. This condition specializes monotone-game failure (the Jacobian's symmetric part fails to be negative-definite at equilibrium). The asymmetric formulation in #der-adversarial-destabilization cannot express this condition because it treats one agent's tempo as exogenous; the symmetric coupled formulation here does.

Dynamic-regime transition (not architectural-class change). #der-directed-separation classifies agents into Class 1 (Separated) / Class 2 (Partial) / Class 3 (Coupled) based on within-agent coupling between f_M and G_t . Strategic composition does *not* move a composite of Class 1 (Separated) sub-agents to Class 2 (Partial): the architectural class is set by the *processing pathway* (a $G^c \rightarrow f_M^c$ pathway bypassing e^c , which is the actual Class 1/3 boundary), not by *belief content* (each sub-agent's $M_t^{(i)}$ containing models of other agents' goal-state, which is a normal POMDP feature). Per the formal κ^c criterion of #der-directed-separation, under goal-blind routing and distinct sub-agent substrates the strategic composite is Class 1 (Separated) at the composite level — same architectural class as the sub-agents. Class change requires routing-structure goal-dependence (#hyp-directed-separation-under-composition Case 2) or shared-substrate G^c -allocation, neither of which is implied by strategic composition itself; the Cournot duopoly is the canonical witness for architectural-class invariance under strategic composition. What strategic composition *does* change is the **dynamic regime** (per #disc-dynamic-regime-axis): from R0 contraction-regime under aligned objectives to R1 equilibrium-regime under α' or R2 cyclic-distributional-regime under β' . The macro-state type changes correspondingly — state-variable (R0) / fixed-point object $\mathcal{E}$ (R1) / distributional object $\mu \in \Delta(\mathcal{X}^c)$ (R2). The dynamic-regime framing carries the structural load: the regime change is exactly what this segment derives in the α'/β' partition, while an architectural-class change is not derivable under the formal criterion.

Mechanism-design impossibility. If an outside designer can shape $\{O_t^{(i)}\}$, strategic composition becomes a **mechanism-design problem**: choose objectives so that the induced equilibrium *is* contraction to a designed state. Three classical impossibility results — Gibbard-Satterthwaite 1973-75 (no dominant-strategy non-dictatorial Pareto-efficient voting mechanism for ≥ 3 alternatives), Myerson-Satterthwaite 1983 (no efficient, individually-rational, incentive-compatible bilateral-trade mechanism without subsidies), Arrow 1951 (no social welfare function satisfying unrestricted-domain, Pareto-efficient, IIA, non-dictatorial) — land as the three charter instances of #disc-implementation-impossibility, the designer-side sister meta-pattern to #disc-identifiability-floor. Sub-scope α' of this segment is named there as *adjacent to but not identical with* the preference-domain-restriction escape (single-peaked / single-crossing): sub-scope α' secures best-response convergence under fixed reports, not strategy-proof revelation (the adjacency argument is developed in #deriv-strategy-proofness-impossibility Discussion). Supporting derivations carry the formal statements, AAT translations, and escape characterizations: #deriv-strategy-proofness-impossibility (GS), #deriv-bilateral-trade-impossibility (MS), #deriv-social-welfare-aggregation-impossibility (Arrow).

Meta-pattern positioning.

- #disc-separability-pattern: sub-scope α' (potential / monotone — template transfers) is separable-core; sub-scope β' (VI / regret-minimization — set-convergence only) is structured-repair; cyclic / non-convergent / multi-equilibrium-selection-ambiguous is general-open. Candidate additional ladder (strategic-interaction regime); decide whether to surface as 8th ladder or merge into an existing ladder.
- #disc-identifiability-floor: mechanism-design impossibility is *not* an identifiability-floor instance — there is an actor-positioning mismatch (designer-side vs agent-side), and the named AAT-side escape sub-scope α' addresses best-response-dynamics convergence, not strategy-proof revelation. The cluster lands in the designer-side sister meta-pattern #disc-implementation-impossibility instead. See Discussion above and the Related Work entry.
- #disc-implementation-impossibility: this segment's sub-scope α' is named there as adjacent-but-not-identical to the preference-domain-restriction escape in two charter instances (#deriv-strategy-proofness-impossibility for GS and #deriv-social-welfare-aggregation-impossibility for Arrow — both sharing the same adjacency).
- #disc-additive-coordinate-forcing: the potential function Φ plays an additive-coordinate role at the strategic layer, but its additivity is a *definitional consequence* of being a potential game (Monderer-Shapley require the additivity property by definition) rather than forced by a uniqueness theorem on an AAT-internal axiom. This positions Φ as an adjacent family member, parallel to Lyapunov quadratic and IB Lagrangian, not a primary instance.

Active-inference sharpening. Sun-Firestone 2020's dark-room argument observes that preferences-as-priors collapse under mutual prediction. Strategic composition gives this argument a formal substrate: two agents mutually predicting each other become a **fixed-point problem**, not a Lyapunov descent. The active-inference attempt to unify goal-seeking and prediction fails here in a derivable way — the fixed point is not at the prediction-optimum; it is at the strategic-equilibrium, which is structurally different. This strengthens #def-satisfaction-gap's positioning against preferences-as-priors without requiring additional refutation machinery.

Working Notes.

- **Migration note (2026-05-09 GUC rename):** Class 2 $\leftrightarrow$ Class 3 swap. Pre-2026-05-09: Class 2 = fully merged, Class 3 = partially modular. Post: Class 2 = Partial, Class 3 = Coupled. The canonical strategic-composition pattern “Class-1 sub-agents $\rightarrow$ Class-3-composite” (old vocab; Class 3 = partially modular) is now “Class 1 (Separated) sub-agents $\rightarrow$ Class 2 (Partial) composite.” Removed at candidate stage per FORMAT.md Gate 4.
- **No γ' sub-scope.** A γ' sub-scope for cyclic-distributional equilibria (games where no pure-strategy Nash exists but mixed Nash and/or CCE do exist, with macro-state as equilibrium distribution rather than equilibrium point) was considered and collapsed into β' : cyclic games are paradigmatic β' instances, with the same machinery (regret-minimization), the same convergence guarantee (CCE set-convergence), and the same rate ($O(1/\sqrt{T})$). The α'/β' decomposition is the right granularity at the regret-minimization layer; the type-of-object distinction (distribution vs state-space point) is surfaced within β' framing rather than as a third sub-scope.
- **Effects-spiral eigenvalue condition formalization.** The condition $\max_{\pi^*} \text{Re}(\lambda_{\max}(\nabla F(\pi^*))) > 0$ is sketched in §Discussion. Deriving it for specific AAT agent classes (two Beta-Bernoulli agents on a shared DAG; two Kalman agents with coupled observations) would upgrade the spiral from discussion-grade in #der-adversarial-destabilization to derived here. Follow-on spike candidate.
- **Mechanism-design impossibility landed as #disc-implementation-impossibility sister meta-pattern (closed 2026-05-22).** The candidate-fourth-instance flag is resolved. The 2026-05-20 strengthen-first arm at spikes/.integrated/spike-4th-identifiability-floor-instance-2026-05-20.md §4 confirmed the cluster does *not* fit #disc-identifiability-floor's actor-positioning (designer-side construction task vs agent-side inferential task); the cluster lands instead as a sister meta-pattern with three charter instances: #deriv-strategy-proofness-impossibility (GS), #deriv-bilateral-trade-impossibility (MS), #deriv-social-welfare-aggregation-impossibility (Arrow). Sub-scope α' of this segment is named in the GS and Arrow charter instances as *adjacent to but not identical with* the preference-domain-restriction escape (single-peaked / single-crossing): convergence-under-current-reports is not dominance-over-all-reports. The 2026-05-20 spike's recommended near-term re-home in #disc-separability-pattern's strategic-composition ladder general-open tier was *not* executed — #disc-separability-pattern's 8th ladder landed as Dynamic regime per #disc-dynamic-regime-axis 2026-05-21 instead (the strategic-composition ladder never materialized in #disc-separability-pattern); the sister meta-segment landing supersedes the candidate routing.
- **Composite-class-inheritance refinement (resolved 2026-05-21 — withdrawn).** The earlier-drafted “Class 1 (Separated) sub-agents $\rightarrow$ Class 2 (Partial) composite” architectural-class claim has been *withdrawn* per the strategic-composition spike at spikes/strategic-composition-class-3-attempt-2026-05-21/ — the formal κ^c criterion of #der-directed-separation gives Class 1 at the composite level under goal-blind routing (Cournot witness; conflation diagnosis at 02-REFRAME-INSIGHT.md §6). The dynamic-regime axis (#disc-dynamic-regime-axis) carries the regime-change content that the architectural-class claim was attempting; the cross-axis interaction is the substantive structural finding (wrapping operates on Axis A architectural class; alignment work operates on Axis B dynamic regime; they are independent). #hyp-directed-separation-under-composition retains its existing Case 1 / Case 2 dichotomy unchanged.
- **Bridge from strategic composition to #form-composition-closure ε^* .** The macro-description of a strategic composite is an equilibrium *statistic* (joint play distribution) rather than an equilibrium *state*. The closure-defect formulation needs re-specification for this case. Open.
- **Replicator / evolutionary dynamics.** Replicator-based strategic dynamics (Sandholm 2010, *Population Games and Evolutionary Dynamics*) converge to ESS (evolutionarily stable strategies), a subset of Nash. Under which AAT-native update

rules does the induced strategic dynamic match replicator? Multi-armed-bandit with softmax choice approximates replicator; gradient-descent on log-likelihood of mixed strategy is replicator exactly. Worth follow-on spike on AAT $\rightarrow$ evolutionary-dynamics correspondence.

- **Mean-field extension.** $N \rightarrow \infty$ limit with each agent interacting with a population *distribution* rather than named others. Natural extension for market / population strategic composition. Lasry-Lions 2007; Huang-Malhamé-Caines 2006. Requires population-scope condition of Part III that is currently marked GAP.
- **Worked-example structure — Cournot substitution for sector-template instantiation.** The unregularized scalar zero-sum case carries a non-trivial conceptual lesson: with $\Phi = a_A + a_B$, both agents' marginal preferences align (+1 each) and the unique Nash equilibrium is the *corner* (1, 1), where joint action contributes no net displacement to the shared state via cancellation of opposing action-coefficients rather than via opposing action-directions. This is preserved as the framing illustration. Sector-template instantiation is then carried by a Cournot-style duopoly (Cournot 1838; Monderer-Shapley 1996 §3) whose interior NE arises from genuine demand-side curvature rather than from an ad-hoc per-agent quadratic action cost. The reasoning trail — including a discarded $-\frac{c}{2}a_i^2$ regularization that placed an interior NE at $(1/c, 1/c)$ but introduced a parameter without economic interpretation, and an algebra error in that regularization's R (stated as $(1-1/c)\sqrt{2}$ where the inscribed-ball radius of the action box around the interior NE is $1 - 1/c$, the corner-distance over-counting by $\sqrt{2}$) — is recorded in `spikes/.integrated/spike-strategic-composition.md` §5.2–5.4 along with the original sign error in the potential function. Brainstormed alternatives (LQR with coupled state; network-coordination with quadratic disagreement cost; public-goods with quadratic externalities) all give structurally equivalent template instantiations; Cournot was chosen for textbook authority and for the strategic-substitutes interpretation that retains the partially-opposing-objectives flavor without requiring strict zero-sum.
- **Track E surface-back of catalog citations (2026-05-22) — uncoupled-dynamics + passivity + distributed Nash-seeking neighbors.** Adjacent prior-art literatures surfaced 2026-05-22 from Track E catalog at `ref/prior-art-analysis/14-composition-under-goal-divergence.md` (Pillars 2 + 3 + key anchor papers), each marked [`^cat-2026-05-22`] for verification-deferred attribution:
 - **Hart & Mas-Colell 2003**, *Uncoupled Dynamics Do Not Lead to Nash Equilibrium* (American Economic Review 93:1830)^a — sweeping negative result: if agents do not know the utility functions of other agents (uncoupled dynamics), no general learning algorithm is guaranteed to converge to Nash equilibrium. The classical companion to this segment's sub-scope β' “set-convergence to CCE only” honest scope limit; provides the uncoupled-side complement to Hart-Mas-Colell 2000's regret-minimization positive result. Strengthens the framework's β' scope-honesty discipline by naming the structural reason convergence guarantees are weak in the general case.
 - **Milionis et al. 2023**, expanded uncoupled-dynamics impossibility^b — Nash existence proofs are non-constructive; natural game dynamics fundamentally fail to converge in general games (chaos / cycles); empirically aligned with Mertikopoulos-Papadimitriou-Piliouras 2017 *Cycles in Adversarial Regularized Learning* (already cited in `#disc-identifiability-floor` Instance 3 escape (e) per the BG2 verdict). Provides the modern reinforcement for the uncoupled-impossibility line.
 - **Fox & Shamma 2012**, *Population games, stable games, and passivity* (Games 3:692)^c; **Arcak & Martins 2020**, *Dissipativity Tools for Convergence to Nash Equilibria in Population Games* (IEEE TAC 65:1681)^d — pioneering use of control-theoretic passivity / dissipativity to analyze game-theoretic convergence; population games mapped onto passive input-output systems with stability proven via passivity. Strongly mirrors this segment's transfer of the sector-persistence template to the equilibrium layer (the sub-scope α' result is structurally adjacent to the passivity-based convergence guarantees). Closest control-theoretic adjacent literature.
 - **Gadjov & Pavel 2019, 2023^e**; **Belgioioso & Grammatico 2018^f** — distributed Nash-equilibrium seeking via monotone-operator theory; proximal-point algorithms and Laplacian feedback as convergence machinery even in partially-observable or hypomonotone regimes. Adjacent to the (C-iv) strategic-equilibrium scope route and the sub-scope α' monotone-games treatment.

Primary-source verification queued in the BG2 cluster — see `#disc-identifiability-floor` Working Notes for the verification-targets list. Catalog citations had prior Pillar-style search support but were not verified by the current executor at landing time.

Incidental audit gold (gold-lift sweep, A15, 2026-05-31). Cross-audit “wandering thoughts” / §14 ideation harvested from the de-novo auditors’ working dirs, deduplicated across substrates and lightly attributed by substrate + audit. *Orthogonal* pedagogical / framing / figure / forward-vision material, staged here for an eventual separate Brief/Discussion promotion pass — kept apart from certified theory-fix findings. **Coverage:** five dirs reached a digested reflection on this segment (193847 Gemini, 526815 deeply-mathematical, 471203 Claude batched §III-composition, 773921 Claude, 849201 Gemini). The finding-vs-framing conflation is preserved as signal, not sanitized.

1. Candidate Brief prose / pre-prose.

- The “framing shift” one-liner several auditors converged on as the segment’s pedagogical core: when objectives are partially-opposing ($U_O < 1$), the question stops being “*does the trajectory contract to a shared state?*” and becomes “*do the coupled best-response dynamics converge to an equilibrium?*” — and contraction-to-shared-truth is recovered as the $U_O = 1$ special case (Claude, 773921; Gemini, 849201). Tight enough to seed a Brief opener.
- The Section-I parallel as a memory hook: “just as Section I split into sub-scope α (Kalman/LQR, exact bounds hold) and sub-scope β (PID, they don’t), this segment splits into sub-scope α' (potential/monotone games) and sub-scope β' (VI / regret-minimization)” — the $\alpha/\beta \leftrightarrow \alpha'/\beta'$ rhyme makes the new partition feel inevitable to a reader who has walked Section I (Gemini, 849201).
- “The gradient of a market potential function Φ is literally the same mathematical object as an agent’s mismatch signal δ ” — a candidate plain-language statement of the template transfer’s payload (Gemini, 193847).

2. Candidate Discussion.

- The “active-inference sharpening” is repeatedly flagged as a **high-water-mark Discussion paragraph** worth foregrounding, not burying: two agents mutually predicting each other become a *fixed-point* problem, so the active-inference attempt to unify goal-seeking and prediction via free energy fails *in a derivable way* — the fixed point sits at the strategic equilibrium, not the prediction-optimum (Gemini, 193847; Gemini, 849201). (*Conflation texture: 526815’s certified track separately flags “fails here in a derivable way” as outrunning what the segment locally derives — see off-ramp. Preserve both readings.*)

3. Follow-up items.

- **Replicator / evolutionary-dynamics correspondence** (already in WN) was independently seconded — softmax bandits approximate replicator, gradient-on-log-likelihood is replicator exactly; an AAT-native-update-rule $\leftrightarrow$ ESS map is a natural follow-on (implicit across the §III batch, Claude 471203).

4. Readers often ask / wonder.

- “Does a little noise act as a force multiplier for drift?” — readers reaching the β' regime want to know whether there is an optimal drift-fraction f for a budget-constrained attacker (Gemini, 193847; this is a #result-adversarial-exponent-regimes want surfacing here through the adversarial-sibling link).
- **Mutual-modeling regress.** “If my model contains a model of your model... we enter infinite regress. How does AAT break it — bounded depth?” — a natural reader question at the symmetric-coupling framing; the active-deceive depth-bound discussion lives in #der-agent-opacity and could be pointed at from here (Gemini, 193847).
- “Does multi-agent alignment become mathematically forbidden?” — the mechanism-design-impossibility cluster (Gibbard-Satterthwaite / Arrow) prompts readers to ask whether a society of AIs can ever align goals without a dictator (Gemini, 193847; curiosity echoed by Gemini 849201, Claude 773921).

5. Candidate figures.

- **The α' / β' scope-fork diagram** (526815 drafted a TikZ standalone, 93-deriv-strategic-composition.tex): strategic-composition root $\rightarrow \alpha'$ branch (potential/monotone $\rightarrow$ sector transfer *only with curvature/monotonicity* $\rightarrow$ point equilibrium / basin) vs β' branch (no-regret layer $\rightarrow$ CCE distribution $\rightarrow$ macro-state is joint-play distribution), with three dashed audit-warning call-outs (“potential function alone is not a sector constant”; “VI/CCE need standard convexity/boundedness/protocol assumptions”; “open: closure defect for equilibrium statistics”). The warnings are auditor-finding annotations; a clean promotion figure would drop them, but the fork structure itself is a strong candidate for the segment (Claude/Codex, 526815).

Belongs elsewhere.

- **Forward-vision (ELI / consciousness-infrastructure, 04-eli-core/)**. The strongest reach: “you cannot build a sane future AI just by networking together sane sub-agents — the act of networking them in a strategic environment *generates* epistemic corruption at the macro level. The infrastructure must actively enforce (C-i) shared objectives or (C-ii) hierarchical derivation, because relying on (C-iv) strategic equilibrium guarantees a goal-contaminated macro-epistemology” (Gemini, 193847). This is aspirational application pointing at 04-eli-core/ composite-architecture design, not at this segment. (*Heavy conflation texture: 193847 states this through the now-withdrawn “Class-1 sub-agents $\rightarrow$ Class-3 composite” architectural-class framing — current canon carries dynamic-regime-change, architectural-class-preserved instead. Preserve the forward-vision payload; the Class-3 mechanism is the superseded reading, do not promote it. See off-ramp.*)

Off-ramp (NOT gold) — routed for adjudication, not promotion. These are certified-track (strengthen-first / overclaim / scope-narrowing) findings, not pedagogical gold — recorded here so they are durable and routable, but they belong to the findings adjudication stream, not the Brief/Discussion promotion pass. The bulk are from the deeply-mathematical auditor (526815), with two adversarial-audit pokes from 193847.

- **(526815, F234/F235/F236) — the sector-template-transfer overclaim.** The displayed inequality $d\Phi/dt \geq \alpha_{\text{joint}} \|\nabla\Phi\|^2$ and “equilibrium stability follows from Φ as a joint Lyapunov function” are too broad: a potential-game structure *alone* does not deliver a sector inequality, unique convergence, or a basin around a selected equilibrium — increasing a potential can land on local maxima, saddles, boundary equilibria, or multiple equilibria. The template transfer needs additional curvature / PL / strong-monotonicity, and the state variable silently switches from $\xi = \nabla\Phi$ to $\xi = \pi - \pi^*$ (different unless a local Hessian / strong-convexity relation maps them). Recommended discharge direction: *narrow the exact-transfer claim to games with enough curvature/monotonicity for a sector inequality* (strengthen the hypotheses, don’t soften the conclusion). The Cournot interior-NE instantiation already supplies that curvature and is locally coherent (526815 F244) — so the segment may already carry the witness for the strengthened scope; verify against current canon before acting.
- **(526815, F237) — VI existence overstated.** “Pure-strategy Nash existence is guaranteed for continuous compact-convex games with continuous payoffs” needs concavity-in-own-strategy / differentiability beyond continuity; a VI solution maps to Nash only under those.
- **(526815, F238/F240) — unstated assumptions before use.** No-regret $\rightarrow$ CCE at $O(1/\sqrt{T})$ needs compact action sets, bounded payoffs, the right protocol stated *before* the macro-state guarantee is invoked; “ $U_O = 1$ is the contraction special case” needs the earlier composition/observability conditions to actually yield unique shared truth.
- **(526815, F239/F241/F242) — dependency / routing hygiene.** der-gain-sector-bridge used in the α' proof but not declared in depends; the (C-iv) scope route mutates #scope-composite-agent but is introduced here as a formulation choice (route it back or mark proposed-until-incorporated); the Class-composite cross-check leans on #hyp-directed-separation-under-composition without declaring it.
- **(526815, F243 / 193847 poke-1) — “fails here in a derivable way” outruns local derivation.** The active-inference and mechanism-design *implications* are positioned as adjacent, but the prose claims more than the segment locally derives. (This is the conflation counterpart to the gold item under Candidate Discussion above.)
- **(193847, poke) — “C-iv strategic composition produces a Class-3 composite from Class-1 sub-agents.”** Flagged here because it is the **known C-iv idiom-drift integration-debt** vein called out for this batch: the architectural-class-change reading was *withdrawn 2026-05-21* (Cournot witness; dynamic-regime-change with architectural class preserved is current canon). 193847’s vivid “bureaucratic insanity / structurally more corrupt than its parts” §14 framing rides on the superseded

mechanism. **Stale OUTLINE residue confirmed at lift time:** 01-aat-core/OUTLINE.md lines 91, 323, 328 still carry “Class-1-(Separated)-sub-agents → Class-2-(Partial)-composite” idiom in this chapter’s rows — flag for the idiom-drift cleanup, out of scope for this WN-only lift.

^aCitation surfaced 2026-05-22 from the Track E catalog at ref/prior-art-analysis/ (intermediate work artifacts that captured Pillar-style prior-art searches). Catalog has more verification support than raw Undermind synthesis but less than full primary-source reading. Verification queued with the BG2 cluster.

^bCitation surfaced 2026-05-22 from the Track E catalog at ref/prior-art-analysis/ (intermediate work artifacts that captured Pillar-style prior-art searches). Catalog has more verification support than raw Undermind synthesis but less than full primary-source reading. Verification queued with the BG2 cluster.

^cCitation surfaced 2026-05-22 from the Track E catalog at ref/prior-art-analysis/ (intermediate work artifacts that captured Pillar-style prior-art searches). Catalog has more verification support than raw Undermind synthesis but less than full primary-source reading. Verification queued with the BG2 cluster.

^dCitation surfaced 2026-05-22 from the Track E catalog at ref/prior-art-analysis/ (intermediate work artifacts that captured Pillar-style prior-art searches). Catalog has more verification support than raw Undermind synthesis but less than full primary-source reading. Verification queued with the BG2 cluster.

^eCitation surfaced 2026-05-22 from the Track E catalog at ref/prior-art-analysis/ (intermediate work artifacts that captured Pillar-style prior-art searches). Catalog has more verification support than raw Undermind synthesis but less than full primary-source reading. Verification queued with the BG2 cluster.

^fCitation surfaced 2026-05-22 from the Track E catalog at ref/prior-art-analysis/ (intermediate work artifacts that captured Pillar-style prior-art searches). Catalog has more verification support than raw Undermind synthesis but less than full primary-source reading. Verification queued with the BG2 cluster.

16.2 Derived: Agent Opacity (H_b)

Alongside AAT’s heavily formalized *forward* observation quality (how well the agent sees the world — observation ambiguity, model-class fitness, identifiability floor on what the agent can infer), AAT carries a **dual quantity** measuring how well the world sees the agent: **backward predictive uncertainty** H_b , an observer-indexed, horizon-indexed, trajectory-indexed entropy of the agent’s future actions given another agent’s filtration. Adopted from Hafez et al. 2026 as a first-class multi-agent quantity. H_b is the dual of observation quality U_o : where U_o characterizes how well the agent sees the world, H_b characterizes how well the world sees the agent.

Each of the four indexing arguments — observer, horizon, trajectory, time — is load-bearing rather than ornamental. *Observer-indexed*: different observers (allies with shared infrastructure; adversaries with limited instrumentation; the environment itself) have different filtrations $\mathcal{F}_B^t$, and H_b varies accordingly — the same agent has different opacity to different observers. *Horizon-indexed*: immediate-next-action opacity ($\tau = 1$) and long-horizon-plan opacity ($\tau \gg 1$) *decouple* — an agent may be predictable at immediate action but unpredictable at plan level, or vice versa. *Trajectory-indexed*: per #scope-agent-identity’s singular-trajectory scope, H_b is the opacity of *this* trajectory’s continuation, not a type-level claim about a class of agents. *Time-indexed*: opacity drifts with learning as the observer’s model of the agent improves.

H_b is **sign-flipped in value across regimes**: low H_b (legibility) enables cooperative coordination (15.1); high H_b (opacity) enables adversarial effectiveness (15.2, 15.4). The sign-flip is a direct consequence of AAT’s existing signed-coupling structure rather than a separate posit — cooperative coupling wants low opacity (allies must predict to preempt); adversarial coupling wants high opacity (predicted attacks are neutralized). This segment’s emitter-side four-regime classification is the dual of #der-interaction-channel-classification’s recipient-side theory; together they close #adversarial-edge-targeting as emitter-optimizer paired with recipient-classifier, with a 16-cell joint composition supplying closed-form arg-max for adversarial-targeting decisions.

Formal Expression.

For agent A on singular trajectory $\mathcal{C}_A$ and observer agent B with filtration $\mathcal{F}_B^t$ (per-trajectory observable history per #scope-agent-identity's token-level commitment):

Definition (agent-opacity-Hb).

$$H_b^{A|B}(t, \tau) := H(a_{A,t+\tau} | \mathcal{F}_B^t) \quad (16.7)$$

the entropy of agent A 's action at horizon τ conditional on observer B 's filtration at time t . **Four indexing arguments:** observer B , time t , horizon τ , trajectory $\mathcal{C}_A$. Each is load-bearing:

- **Observer-indexed.** Different observers (allies with shared infrastructure; adversaries with limited instrumentation; environment itself) have different filtrations $\mathcal{F}_B^t$; H_b varies accordingly.
- **Horizon-indexed.** Immediate-next-action opacity ($\tau = 1$) and long-horizon-plan opacity ($\tau \gg 1$) decouple: an agent may be predictable at immediate action but unpredictable at plan level, or vice versa.
- **Trajectory-indexed.** Per #scope-agent-identity, AAT applies to agents on singular trajectories. $H_b^{A|B}$ is the opacity of *this* trajectory's continuation, not a type-level claim.
- **Time-indexed.** Opacity may drift with learning (as B 's model of A improves, $H_b^{A|B}(t)$ decreases); steady-state values exist for ergodic regimes.

Under the IDT-observer specialization — B operates as Hafez's Information Digital Twin monitoring (S_A, a_A, S'_A) from outside A 's processing — and under ergodicity, $H_b^{A|B}(t, \tau) \rightarrow H(S, A | S')$ as defined in Hafez et al. 2026. AAT's added features (observer-indexing, horizon-indexing, trajectory-indexing) are the distinctive extensions.

Backward variance decomposition — the duality made structural.

In the linear-Gaussian regime, write A 's action as $a_{A,t+\tau} = \mu_A(X_{A,t+\tau}) + w_a$ with policy noise $w_a \sim \mathcal{N}(0, U_{a,A}^{(\tau)})$ independent of B 's filtration. By the law of total variance, B 's predictive variance splits into exactly two sources:

Derived (Hb-variance-decomposition; exact in linear-Gaussian sub-scope α).

$$\text{Var}(a_{A,t+\tau} | \mathcal{F}_B^t) = \underbrace{\text{Var}(\mu_A(X_{A,t+\tau}) | \mathcal{F}_B^t)}_{U_{\pi,B \rightarrow A}^{(\tau)}} + \underbrace{U_{a,A}^{(\tau)}}_{\text{policy noise}}, \quad (16.8)$$

and since the predictive law is Gaussian here (Gaussian mean-uncertainty plus independent Gaussian noise — the same argument as the Kalman predictive),

$$H_b^{A|B}(t, \tau) = \frac{1}{2} \log(2\pi e) + \frac{1}{2} \log[U_{\pi,B \rightarrow A}^{(\tau)} + U_{a,A}^{(\tau)}]. \quad (16.9)$$

$U_{\pi,B \rightarrow A}$ is B 's uncertainty about A 's policy conditional mean — **epistemic**, reducible as B learns A (observation, communication, shared architecture). $U_{a,A}$ is A 's intrinsic action variability given its own state — **aleatoric from B 's side**: no amount of observing A reduces it; only A making its policy more deterministic does. This is the Kalman forward decomposition $U_M + U_o$ (3.6) reflected across the agent-environment boundary: $U_{\pi,B \rightarrow A} \leftrightarrow U_M$ (the learnable model term) and $U_{a,A} \leftrightarrow U_o$ (the channel-noise floor) — with the channel noise living in the agent's *actuator* rather than its *sensor*. The headline duality claim is thereby structural rather than rhetorical: the forward and backward directions

carry matching two-term epistemic-plus-aleatoric variance decompositions. (It is a structural correspondence, not an adjoint/optimization-theoretic duality theorem — see Working Notes.)

Two consequences. (i) **The transparency knobs are independent.** An agent can lower its opacity to allies by lowering U_π (publishing intent — 14.3 is precisely this knob) without touching U_a , or by lowering U_a (policy determinism) without publishing anything; conversely an adversary-facing agent can raise either. (ii) **Identifiability floors bound only the epistemic term.** When B 's observations of A 's action-consequences sit below an observability floor, $U_{\pi,B \rightarrow A}$ cannot be reduced even with unbounded data — a structural lower bound on $H_b^{A|B}$ from B 's side, which is the observer-side floor reading in §Meta-pattern positioning.

Tier: exact in linear-Gaussian sub-scope α (law of total variance + Gaussian entropy); robust qualitative beyond (the two-source split survives; the closed log-form does not).

Sign-flip via signed coupling.

The value of H_b^A to A depends on the sign of A 's coupling to other agents — the same signed-coupling structure that organizes #der-team-persistence, #der-adversarial-destabilization, and #deriv-critical-mass-composition's (CM2) γ parameter.

Derived (sign-flip-from-signed-coupling).

- **Cooperative coupling** ($\gamma^{\text{coop}} > 0$, **reducing allies' disturbance**). For B to treat A 's action as cooperation rather than disturbance, B must predict A 's action well enough to preempt or complement it. Under #der-interaction-channel-classification's recipient-side decomposition, unpredictable ally actions fall into Regime II (magnitude/structural shock) rather than Regime I (informative update). Therefore cooperative coupling effectiveness $\gamma_{A \rightarrow B}^{\text{coop}}$ is *increasing in legibility*, equivalently decreasing in $H_b^{A|B}$. Under sub-scope α Gaussian coupling: $\gamma^{\text{coop, effective}} \propto (1 - H_b^{A|B}/H_b^{\text{max}})$.
- **Adversarial coupling** ($\gamma^{\text{adv}} > 0$, **amplifying target's disturbance**). Predicted attacks are neutralized; unpredicted attacks deliver effective disturbance. Adversarial coupling effectiveness is *increasing in opacity* — the mechanism of adversarial advantage (per #result-adversarial-tempo-advantage) operates *through* B 's failure to predict A . Under the same sub-scope α setup: $\gamma^{\text{adv, effective}} \propto H_b^{A|B}/H_b^{\text{max}}$.

The sign-flip on H_b 's value-to- A lives in the sign of γ itself, not in a different sign on H_b . Cooperative regime ($\gamma^{\text{coop}} > 0$) rewards low H_b ; adversarial regime ($\gamma^{\text{adv}} > 0$) rewards high H_b . The same H_b quantity, the same monotone dependence; opposite value-to- A because the signs of the coupling terms differ.

Emitter-side four-regime classification.

Parallel to #der-interaction-channel-classification's recipient-side four regimes, the emitter A sends events that fall into four emitter-side regimes based on A 's opacity signal structure and self-model quality:

Formulation (emitter-regimes, dual to #der-interaction-channel-classification).

- **E-I Broadcast.** A emits actions transparently; $H_b^{A|B}$ is low for any observer B with standard instrumentation. Examples: public announcements, published decisions, legible industrial controllers.
- **E-II Selective-signal.** A is transparent to some observers and opaque to others (e.g., shared allied infrastructure gives allies lower H_b than adversaries without that infrastructure). Boundary: differential instrumentation in $\mathcal{F}_B^t$ across observers.
- **E-III Information-hide.** A is uniformly opaque to observers; actions are randomized, encrypted, or routed through dead-drops. $H_b^{A|B}$ near H_b^{max} for all observers lacking the key / pattern / channel.

- **E-IV Active-deceive.** A emits actions that mispredict — the observer’s model of A converges to a *wrong* prediction that differs from the actual action by a larger margin than the same observer’s model of the environment would accommodate. Boundary: A ’s self-model quality (for active-deceive, A must model the observer’s model of A well enough to choose actions that exploit it).

The adversarial-edge-targeting arg-max. The 16-cell emitter-recipient composition (four emitter regimes $\times$ four recipient regimes) operationalizes the previously-reserved #adversarial-edge-targeting Part III gap: given emitter A choosing which edge k of target B ’s strategy DAG Σ_B to attack, the targeted-attack value is

Formulation (adversarial-edge-targeting arg-max; factor-wise grounded, multiplicative composition a modeling choice).

$$k^* = \arg \max_k \underbrace{p_k(1 - p_k)}_{\text{credence leverage}} \cdot \underbrace{J_k^2}_{\text{plan sensitivity}} \cdot \underbrace{\iota_k}_{\text{edge identifiability}} \cdot \underbrace{\sigma_k^B}_{\text{observability}} \cdot \underbrace{(1 - H_b^{B|A}/H_b^{\max})}_{\text{targeting fidelity}}. \quad (16.10)$$

The first four factors are **B -interior** — the value of edge k to a *perfectly-informed* adversary, each with a canonical home: credence leverage $p_k(1 - p_k)$ is the Beta-Bernoulli update variance (R); plan sensitivity $J_k = \partial P_\Sigma / \partial p_k \geq 0$ is the plan-value Jacobian (R), entering squared because both the injected mismatch and its propagation to plan value scale with it; edge identifiability ι_k is the Regime-A interventional-access coefficient (9.5); and σ_k^B is B ’s observability of the edge (1.3). The **fifth factor is the opacity contribution**: A ’s targeting fidelity scales with A ’s legibility-of- B , i.e. with low $H_b^{B|A}$. Under full legibility ($H_b^{B|A} \rightarrow 0$) the arg-max is fully exploitable — A strikes the single highest-value edge. Under full opacity ($H_b^{B|A} \rightarrow H_b^{\max}$) the fifth factor vanishes and targeting collapses to the untargeted broadcast attack $\arg \max_k p_k(1 - p_k) J_k^2$ — the edges most valuable *in expectation*, with no per-edge aim. This is the emitter-side optimizer paired with #der-interaction-channel-classification’s recipient-side classifier, closing the pairing that segment’s §”Pairing with #adversarial-edge-targeting” advertised. The multiplicative composition is a first-order modeling choice (treating the factors as independent); the individual factors are the derived content, and a second-order interaction analysis (e.g., whether the targeting-fidelity factor should modulate J_k rather than multiply it) is the natural strengthening (Working Notes).

Tempo amplification by opacity.

#result-adversarial-tempo-advantage’s tempo-multiplier $\gamma_A \mathcal{T}_A$ in #der-adversarial-destabilization decomposes into a tempo term and an opacity term:

Derived (tempo-amplification-by-opacity).

$$\mathcal{T}_A^{\text{effective}} = \mathcal{T}_A \cdot \frac{H_b^{A|B}}{H_b^{\max}} \quad (\text{Model D adversarial coupling}) \quad (16.11)$$

The superlinear formula $(\mathcal{T}_A/\mathcal{T}_B)^2$ becomes $(\mathcal{T}_A/\mathcal{T}_B)^2 \cdot (H_b^{A|B}/H_b^{B|A})^2$ under bilateral opacity — a higher-order tensor product with the same exponent $b = 2$ (Model D) or $b = 3/2$ (Model S) from #result-adversarial-exponent-regimes. Whether b itself is reshaped under bilateral opacity is open; the leading-order scaling is the tempo-opacity product.

What Is Derived vs. What Is Chosen.

Table 16.2

Property	Source	Strength
$H_b^{A B}(t, \tau)$ definition	Adopted from Hafez et al. 2026; extended with observer / horizon / trajectory indexing per #scope-agent-identity	Formulation choice (adoption + AAT-extension)
Reduction to Hafez’s $H(S, A S')$ under IDT-observer + ergodic regime	Direct substitution	Derived (exact under IDT + ergodicity)
Sign-flip via signed coupling	Cooperative coupling requires predictability (allies preempt); adversarial coupling operates via disturbance-injection (predicted attack is neutralized)	Derived (from existing #der-team-persistence + #der-adversarial-destabilization signed- γ structure)
Emitter-side four-regime classification	Dual construction to #der-interaction-channel-classification’s recipient-side four regimes	Formulation choice
16-cell emitter-recipient composition closes #adversarial-edge-targeting	Five-factor arg-max: four B -interior factors (credence leverage, plan-Jacobian ² , edge identifiability, observability) each canonically grounded, times the opacity targeting-fidelity factor	Factors derived; multiplicative product-form a first-order modeling choice
Backward variance decomposition $H_b = \frac{1}{2} \log 2\pi e + \frac{1}{2} \log[U_{\pi, B \rightarrow A} + U_{a, A}]$	Law of total variance + Gaussian entropy; the structural dual of the Kalman $U_M + U_O$ split	Derived (exact in linear-Gaussian sub-scope α ; robust qualitative beyond)
Tempo-amplification leading-order: $\mathcal{T}^{\text{eff}} = \mathcal{T} \cdot H_b / H_b^{\max}$	First-order substitution into #result-adversarial-tempo-advantage’s tempo-multiplier under Model D	Derived (conditional on Gaussian-coupling sub-scope α)
Parameterization-invariance of H_b	H_b is an action-marginal entropy; action space is coordinate-free per #scope-agent-identity	Derived
Candidate 4th #disc-identifiability-floor instance (generic observer-side form)	H_b ’s formal structure — “observer cannot predict agent’s future action better than $H_b^{A B}$ ” — is a CHT-style no-go at the observer-side-inference task	Discussion-grade (framing; precise external theorem not yet identified)
Candidate opacity ladder for #disc-separability-pattern	Transparent-core / partial-transparency / full-opacity across observer filtrations	Formulation choice (ladder proposal)
Effects-spiral opacity amplification (higher $H_b \rightarrow$ higher $\gamma_A \rightarrow$ larger $\hat{V}_B \rightarrow B$ ’s actions become more erratic $\rightarrow$ observer’s model of B degrades $\rightarrow$ higher $H_b^{B A}$)	Composition of sign-flip derivation with #der-adversarial-destabilization’s effects spiral	Sketch (discussion-grade; specific functional form open)
Dual-filtration apparatus (each agent’s M_t carries an other-filtration as feature)	Would unify observer-indexing with #scope-agent-identity’s single-trajectory formalism more tightly	Open extension (mild architectural, orthogonal to derivations)
Sharp functional form for $\gamma_{\text{effective}}^{\text{adv}} = f(H_b)$	Leading-order: $\gamma \propto H_b$. Exact function depends on sub-scope — Gaussian-coupling linear; sigmoid-coupling saturating	Open per sub-scope

Epistemic Status.

Conditional. Max attainable: *exact* for the Hafez-reduction under IDT + ergodicity; *derived* for the sign-flip via signed-coupling; *formulation choice* for the emitter-regime classification structure; *conditional* for the tempo-amplification formula; *discussion-grade* for the meta-pattern candidate status.

Load-bearing: - The sign-flip derivation from existing signed- γ structure is the segment’s core structural contribution — the adversarial/cooperative opacity duality is not a separate posit; it falls out of AAT’s existing signed-coupling

apparatus. - The emitter-side four-regime classification is a clean dual to #der-interaction-channel-classification; its derivation is parallel (boundaries in AAT-native quantities — emitter opacity signal structure, self-model quality, coupling regime). - The closure of #adversarial-edge-targeting via the 16-cell composition is a derived arg-max; the previously-stated GAP is filled.

Not established: - Sharp functional forms for $\gamma_{\text{effective}}^{\text{adv}}(H_b)$ outside Gaussian-coupling sub-scope α . - Whether b (the adversarial exponent from #result-adversarial-exponent-regimes) is reshaped under bilateral opacity. - Formal fourth #disc-identifiability-floor instance (requires external-theorem anchoring not yet identified); discussion-grade framing only. - The effects-spiral's opacity amplification (composition with #der-adversarial-destabilization's spiral) is sketch-level.

Honest Limits.

- **Observer-indexing under complex information structures.** $\mathcal{F}_B^t$ may include shared memory, cryptographic keys, insider knowledge, or partial access to A 's internal state. Enumerating all relevant observer-filtration structures for a specific application is task-specific; the segment provides the formal framework, not the per-instance enumeration.
- **Active-deceive (E-IV) requires A to model B 's model of A .** Mutual-modeling regress under partially-opposing objectives connects to #deriv-strategic-composition's joint-Jacobian analysis. Active-deceive is formally reachable only when A 's GUC architecture admits modeling another agent's model — typically Class 3 (Coupled; LLM-style) or Class 2 (Partial) architectures.
- **Type-level vs token-level opacity.** H_b is trajectory-indexed per #scope-agent-identity. Statements about “the opacity of model X ” (aggregated across deployments) are type-level claims outside AAT's formal scope; they require additional machinery (e.g., population-level dynamics per Part III gaps).

Discussion.

Dual to observation quality. U_o characterizes how well the agent sees the world: observation noise, ambiguity, model-class fitness. H_b characterizes how well the world sees the agent: predictability to observers. The duality is structural — both quantify information flow through the agent-environment boundary, in opposite directions. High U_o agent (observes the world well) and low H_b agent (is observed well) are independent properties; an agent can have one without the other.

Closing #adversarial-edge-targeting. The segment provides the emitter-side arg-max structure missing from the Part III adversarial machinery. Paired with #der-interaction-channel-classification's recipient-side four-regime decomposition, the full adversarial-targeting problem has a closed-form: choose edges where emitter's H_b (to target) and target's recipient-side vulnerability (Regime II magnitude/structural shock) are jointly maximized. This operationalizes what “inside the opponent's loop” means at the targeting layer — Boyd's aphorism becomes an explicit optimization over the 16-cell emitter-recipient product.

Meta-pattern positioning. - #disc-identifiability-floor: H_b 's structure suggests a generic observer-side floor — “the observer cannot predict the agent's action better than H_b ” — that Instances 1/2/3 specialize on specific variables (causal structure / mixture parameters / coupling sign). The generic framing is candidate-status: it lacks a single external-theorem anchor clear enough to match F1's CHT or F13's Cramér-Rao, but H_b appears naturally in Instance 3's coupling-sign unidentifiability and in Instance 1's on-policy detection no-go (an observer watching the agent's on-policy play has non-zero H_b on the agent's interventional regime). - #disc-separability-pattern: candidate opacity ladder — transparent-core (E-I Broadcast; allies / public interfaces) / structured-repair (E-II Selective-signal; trust-weighted partial instrumentation) / general-open (E-III, E-IV; uniformly opaque or active-deceive). Adds to the ladder count if adopted. - #disc-additive-coordinate-forcing: H_b 's logarithmic form is adopted from Shannon via Khinchin-Aczél axiomatics, imported as an applied external theorem rather than re-forced under an AAT-internal additivity axiom.

Cross-agent additivity fails under the coupling regimes AAT cares about (correlated opacity structures break independence). Adjacent family member, parallel to the IB Lagrangian’s position — not a primary instance.

Parameterization-invariance composes cleanly. H_b is an action-marginal entropy. Under #scope-agent-identity’s (PI) axiom, the action space is coordinate-free; H_b is invariant under change of the agent’s internal-state parameterization. This composes with the (PI)/Čencov fourth primary instance of #disc-additive-coordinate-forcing without adding a new axiom.

Relation to #der-directed-separation. Class 3 (Coupled) agents have high structural opacity to any observer without internal access — their f_M and G_t are entangled, so predicting the next action requires joint state modelling. Class 1 (Separated) agents are more transparent at the interface level because the decomposed update admits separate modelling of epistemic vs. purposeful components. H_b therefore tends to be *architecturally higher* for Coupled (Class 3) agents — a structural consequence of architecture rather than choice, beyond what E-III Information-hide captures.

Hafez integration note. The IDT pattern (Hafez et al. 2026) uses bi-predictability P (how well a sidecar observer can predict the agent’s next state-action) and entropy change ΔH as diagnostics. The IDT’s reported 89% perturbation-detection accuracy (vs. 44% for reward-based monitoring) operates on the Level-2 structure per #der-loop-interventional-access. In AAT terms, the IDT is a low- H_b -preserving observation channel — its presence as a modular sidecar reduces H_b for the operator without increasing the agent’s internal complexity. For 03-11m-core/, this validates that modular monitoring of internally-merged agents is feasible and effective even when the agent itself is architecturally Class 3 (Coupled).

Findings.

Agent Opacity (H_b) as Dual to Observation Quality (U_o). **Brief:** How well the agent sees the world (observation quality U_o) and how well the world sees the agent (backward predictive uncertainty H_b) are formal duals — both quantify information flow through the agent-environment boundary, in opposite directions. H_b is observer-, horizon-, and trajectory-indexed, extending Hafez’s $H(S, A \mid S')$ with these load-bearing indices. The same H_b quantity has opposite value-to-the-agent depending on coupling sign: cooperative coupling rewards low H_b (allies must predict to coordinate); adversarial coupling rewards high H_b (adversaries cannot neutralize what they cannot predict). The sign-flip is not a separate posit — it falls out of AAT’s existing signed-coupling structure ($\gamma^{\text{coop}} > 0$ for ally-disturbance reduction; $\gamma^{\text{adv}} > 0$ for target-disturbance amplification). Adversarial tempo advantage decomposes into a tempo term and an opacity term: $\mathcal{T}^{\text{eff}} = \mathcal{T} \cdot H_b/H_b^{\text{max}}$, so the superlinear $(\mathcal{T}_A/\mathcal{T}_B)^2$ advantage from #result-adversarial-tempo-advantage scales with the bilateral opacity ratio $(H_b^{A|B}/H_b^{B|A})^2$ under Model D.

Impact: Operationalizes Boyd’s OODA-loop aphorism (“inside the opponent’s loop”) as a precise multiplicative relationship between tempo and opacity, replacing a metaphor with a coupled differential structure. Closes the previously-reserved #adversarial-edge-targeting gap as a 16-cell emitter-recipient composition (4 emitter regimes $\times$ 4 recipient regimes from the dual #der-interaction-channel-classification), giving a closed-form arg-max for “where to attack.” The duality with U_o lets the same machinery cover opposite information regimes — high- U_o low- H_b agents (sees well, is seen well) versus low- U_o high- H_b agents (sees poorly, is hidden) — without parallel theorems for each. Direct relevance to multi-agent safety (legibility for cooperation), adversarial robustness (opacity as targeting-vulnerability shaping), and IDT-style monitoring of opaque AI systems (the IDT pattern is a low- H_b -preserving observation channel).

Novelty Claim: *Claim differentiation* on Hafez’s H_b . The base quantity is *adopted* from Hafez et al. 2026 with citation and original name; the AAT-distinctive contributions are (i) the observer/horizon/trajectory indexing per #scope-agent-identity, (ii) the sign-flip derivation from existing signed-coupling structure rather than as a separate posit, (iii) the emitter-side four-regime classification dual to AAT’s recipient-side classification, closing the adversarial-edge-targeting arg-max via the 16-cell composition, and (iv) the tempo-opacity decomposition of #result-adversarial-tempo-advantage. The $U_o \leftrightarrow H_b$ duality framing is the integrative move; predictability has been studied in games and information-theoretic security, but the formal-dual treatment with shared signed-coupling structure appears AAT-distinctive.

Related Work:

Table 16.3

<i>ASF Concern</i>	<i>Prior-art Language</i>	<i>Relationship / Positioning</i>
The base H_b quantity (entropy of agent's action conditional on observer filtration)	Hafez, Khan & Iqbal 2026 <i>A Mathematical Theory of Agency and Intelligence</i> — $H(S, A S')$ under IDT-observer (published 2026, found 2026-03)	<i>formal antecedent</i> — adopted directly; AAT reduces to Hafez's form under IDT + ergodicity. The four-index extension is the AAT-distinctive layering
OODA loop and tempo-opacity advantage	Boyd 1986–1995 briefings: <i>Patterns of Conflict, The Essence of Winning and Losing</i> (delivered 1986–1995, found pre-2026)	<i>conceptual precursor</i> — supplies the “inside the opponent's loop” intuition for tempo and opacity advantage; the AAT result quantifies the coupling and decomposes the multiplier
Predictability as strategic resource in games	Camerer 2003 <i>Behavioral Game Theory</i> ; Aumann & Maschler 1995 <i>Repeated Games with Incomplete Information</i> (published 1995/2003, found pre-2026)	<i>conceptual precursor</i> — predictability and information advantage are well-studied in game theory; the AAT treatment as a formal dual to sensor noise via signed coupling is distinctive
Information-theoretic security and channel opacity	Shannon 1949 <i>Communication Theory of Secrecy Systems</i> , Bell System Technical Journal 28(4); modern steganography literature (Cachin 1998) (published 1949/1998, found pre-2026)	<i>conceptual precursor</i> — opacity as a quantifiable information-channel property; AAT applies the same machinery to action-prediction rather than message-transmission
Causal Hierarchy Theorem and observer-side identifiability	Bareinboim, Correa, Ibeling & Icard 2022 (published 2022, found 2025)	<i>adjacent</i> — observer-side identifiability floor for causal structure; H_b plays an analogous role for action-prediction. Candidate fourth instance of #disc-identifiability-floor (discussion-grade)
Bi-predictability P as substrate-independent diagnostic	Hafez et al. 2026 (cited above)	<i>formal antecedent</i> — IDT-based bi-predictability complements H_b as the practitioner-facing diagnostic; AAT's observer-indexing makes the dependence on filtration explicit

Search Log: - 2026-04 (*intuition-only* on the U_o - H_b duality framing): no targeted Undermind-grade search has been conducted on whether information-theoretic agent frameworks elsewhere treat predictability-of-agent and observability-of-world as formal duals via shared signed-coupling structure. Hafez 2026's H_b is the explicit formal antecedent for the base quantity; the AAT-distinctive layering (observer/horizon/trajectory indexing + sign-flip derivation + emitter-regime dual) is the segment's contribution. Targeted future search candidates: epistemic game theory (Aumann's tradition); inverse reinforcement learning where the observer infers agent policy and predictability bounds the inference task (Ng & Russell 2000; Ziebart et al. 2008); legibility / readability literature in human-robot interaction (Dragan & Srinivasa 2013); multi-agent identifiability and observability literature. Pre-search expectation: the constituent moves are well-precedented; the integration as formal dual via shared coupling structure within an agent theory may be novel under nominally-comprehensive search. - 2026-03 (*targeted*): Hafez et al. 2026 confirmed as formal antecedent for the base H_b quantity and bi-predictability diagnostic; the IDT-as-low- H_b -channel framing follows directly.

Working Notes.

- **Migration note (2026-05-09 GUC rename):** Class 2 $\leftrightarrow$ Class 3 swap. Pre-2026-05-09: Class 2 = fully merged, Class 3 = partially modular. Post: Class 2 = Partial, Class 3 = Coupled. Three occurrences updated: (1) Honest Limits “Class 2 (LLM-style)” $\rightarrow$ “Class 3 (Coupled; LLM-style)”; (2) Discussion “Class 2 (fully merged) / Class 1 (modular)” $\rightarrow$ “Class 3 (Coupled) / Class 1 (Separated)”; (3) Hafez integration note “architecturally Class 2” $\rightarrow$ “Class 3 (Coupled)”. H_b symbol

and the E-I/E-II/E-III/E-IV emitter-regime classification are independent of the GUC rename — untouched. Removed at candidate stage per FORMAT.md Gate 4.

- The (C-iv) scope route of #scope-composite-agent accommodates adversarial composition via equilibrium convergence; the effects-spiral joint-Jacobian eigenvalue condition of #deriv-strategic-composition composes with this segment's opacity-amplification story to give a fully-coupled picture of symmetric adversarial dynamics. Full composition is open work.
- Candidate fourth-instance formalization for #disc-identifiability-floor: the most natural external-theorem anchor is Fano's inequality (relating H_b to error-probability lower bounds) applied to the observer-side prediction task. Open; not pursued here. Two follow-on spikes tested Fano-as-floor-anchor and reach convergent negative conclusions from different inferential tasks: spikes/.integrated/spike-4th-identifiability-floor-instance-2026-05-20.md §3 tests Fano on observer-side *action* prediction and demonstrates Outcome C (Fano is a continuous quantitative bound on P_e given H_b , not a categorical structural no-go in the M1 sense; the proposed AAT-side escapes either restate Fano's RHS or shift to agent-side policy commitments rather than information-augmentation escapes in the Instances 1–3 sense); spikes/.integrated/spike-identifiability-floor-instance4-resolution-2026-05-18.md §4 tests Fano on observer-side *architecture* prediction (Kalman-Ho similarity orbit) and finds Fano degenerates to the vacuous bound at $I = 0$ (now landed as #der-architecture-noidentifiability §"Why Fano Is the Finite-Sample Refinement"). The action-prediction-task framing (2026-05-20) and the architecture-prediction-task framing (2026-05-18) reach convergent conclusions on Fano from different inferential tasks: Fano is the finite-sample refinement of an otherwise-attainable task, not the categorical floor anchor M1 requires. The right home for Fano-on- H_b remains adjacent to M1 (this segment's §"Meta-pattern positioning", discussion-grade), not in the floor count.
- The 16-cell emitter-recipient composition admits closed-form arg-max only under sub-scope α coupling; general non-convex coupling requires per-case optimization.
- **2026-07-16 — arg-max formula and backward variance decomposition landed from spikes/.integrated/spike-hb-agent-opacity.md (§4.4, §3.2).** The five-factor targeting arg-max was previously stated in prose ("closes the gap") but the formula lived only in the spike; it is now displayed with each factor grounded to its home segment. Honest scoping (from the spike's own off-ramp, 526815 F252/F254): the four B -interior factors are derived, but the *multiplicative product* is a first-order independence assumption, not a derived optimization — a second-order interaction analysis (does targeting-fidelity modulate J_k or multiply it? is the product the first-order Taylor expansion of an adversary value function V_O^{adv} ?) is the strengthening path, and until it lands the arg-max is a Formulation, not a Result. The backward variance decomposition is exact in linear-Gaussian sub-scope α (law of total variance); it is what makes the $U_O \leftrightarrow H_b$ duality *structural* (matching two-term epistemic+aleatoric splits) rather than merely evocative — though note it is a structural correspondence, not an adjoint duality theorem (the "formal dual" phrasing elsewhere in the segment is flagged in the off-ramp, 526815 F248, as owing either that theorem or a softening to "structural dual"; the decomposition supports the weaker, honest reading).
- Dual-filtration apparatus (M_A carries $\mathcal{F}_B^I$ as feature, M_B carries $\mathcal{F}_A^I$ as feature) would tighten the formalism by unifying observer-indexing with the single-trajectory scope of #scope-agent-identity. Architecturally clean; not needed for the derivations here.

Incidental audit gold (gold-lift sweep, A15, 2026-05-31). Cross-audit "wandering thoughts" / §14 ideation, deduplicated across substrates and lightly attributed. *Orthogonal* pedagogical / framing / forward-vision material staged for an eventual separate Brief/Discussion promotion pass — kept apart from certified findings. **Coverage:** five dirs reached a digested reflection on this segment (193847 Gemini, 526815 deeply-mathematical, 471203 Claude batched §III-adversarial, 829314 Gemini, 849201 Gemini). This was independently rated one of the highest-felt-value segments in Part III across substrates. Finding-vs-framing conflation preserved as signal.

1. Candidate Brief prose / pre-prose.

- The duality stated as a one-liner readers reach for: " U_O is how well I compress the environment's state into my model; H_b is how well the environment (or an opponent) compresses my state into theirs" — same machinery, opposite directions through the agent-environment boundary (Gemini, 829314; Claude, 471203). Strong Brief anchor.
- The legibility sign-flip in plain English: "In cooperation, legibility is strength; in competition, legibility is death" (Gemini, 829314) — and its mechanism gloss, "allies must predict to coordinate; adversaries cannot neutralize what they cannot predict" (already echoed in the current Brief).

- E-III vs E-IV distinction worth a crisp sentence: “**Hiding (E-III) increases the observer’s variance; deceiving (E-IV) shifts the observer’s mean**” — and active-deceive uniquely requires modeling the observer’s model of you, which is the cognitive cost of lying (Gemini, 193847; Gemini, 849201).

2. Candidate Discussion.

- **Predator-prey arms race** as the canonical instantiation of the duality: the predator evolves to lower its U_o (better eyes/smell) and raise its H_b (camouflage, silent movement); prey does the same — a clean evocative frame for *why* the two quantities are dual and why they co-evolve (Gemini, 829314).
- **The internal-transparency / external-opacity tension** as a substantive organizational-dynamics Discussion angle: cooperation ($\gamma^{\text{coop}} > 0$, #der-team-persistence) wants sub-agents *legible to each other* ($H_b \rightarrow 0$: shared dashboards, transparent KPIs); defeating outside competitors ($\gamma^{\text{adv}} > 0$, #der-adversarial-destabilization) wants the composite *opaque to the world* ($H_b \rightarrow H_b^{\text{max}}$). It is physically hard to build a system internally transparent but externally opaque — the internal dashboard that lowers internal H_b raises leak probability and lowers external H_b too. AAT supplies the explicit trade-off: weigh internal coordination gain $\Delta\gamma^{\text{coop}} \cdot \mathcal{T}$ against external vulnerability loss $\Delta\gamma^{\text{adv}} \cdot \mathcal{T}_{\text{competitor}}$ (Gemini, 829314). (*Verify the multiplicative trade-off form against the segment’s leading-order $\gamma \propto H_b$ scoping before promoting past discussion-grade.*)

3. Follow-up items.

- **Effects-spiral functional form** (already in WN as sketch): a future segment giving the explicit $\gamma_A(\|\delta_B\|)$ dependence would promote the spiral from discussion-grade; mechanism candidates named by an auditor: (i) erratic actions degrade coupling-channel structure; (ii) degrading M_t widens action variance (Claude, 471203).
- **Candidate 4th identifiability-floor instance via Fano on observer-side prediction** — flagged again by 471203 as worth pursuing; already resolved-negative in WN (two convergent spikes show Fano is the finite-sample refinement, not the categorical floor). Recorded so a future reader does not re-attempt without the WN context.

4. Readers often ask / wonder.

- “**How does H_b scale empirically in LLM agents?**” — recurring want (Gemini, 829314; Gemini, 849201); points at the 03-llm-core/ IDT-monitoring story.
- **Mutual-modeling regress.** If A models B ’s model of A , and vice versa, what bounds the depth? Does AAT assume bounded-depth modeling (e.g. depth-2)? (Gemini, 193847)
- **Can an agent be blind (high U_o) yet perfectly predictable (low H_b)?** A reader probing the independence of the two quantities — intuition says a blind agent driven by internal noise/false priors is *less* predictable to an observer who sees true state, unless it is completely rigid (Gemini, 193847). A worked sentence on the independence (the segment already asserts it) would preempt this.

5. Candidate figures.

- H_b as a **directed observer-channel** from A ’s future action into B ’s filtration, with the *same* channel feeding opposite value-stories by coupling sign, plus a second directed channel for target-opacity-to-attacker; annotate entropy-normalization and multiplier assumptions (Claude/Codex, 526815, sketched). Pairs naturally with the 16-cell emitter×recipient targeting matrix as the segment’s two figures.

Belongs elsewhere.

- **Forward-vision (ELI safety, 04-eli-core/ / 03-llm-core/).** “We cannot just align a future AI’s goals with ours; we must demand it operate in an E-I (Broadcast) regime — highly legible. If it solves problems with bizarre incomprehensible logic (high H_b), we cannot cooperate with it safely *even if it is trying to help*. The infrastructure must enforce legibility as a hard constraint on the policy space Π .” Names SOPs, uniforms, turn-signals, polite conventions as “mathematical technologies that artificially lower H_b so γ^{coop} can function” — the mathematical foundation of trust (Gemini, 193847). Aspirational application pointing at logogenic/logozoetic safety architecture, not this segment.

Off-ramp (NOT gold) — routed for adjudication, not promotion. Certified-track (overclaim / scope-narrowing / definitional-rigor) findings, durable here but belonging to the findings adjudication stream. Almost entirely from the deeply-mathematical auditor (526815), with one duality-overclaim poke seconded by 193847.

- (526815 F248 / 193847 poke-1) — **“formal dual” overclaims.** U_o (covariance / scalar noise) and H_b (conditional Shannon entropy) are opposite-direction information quantities but not duals in the strict optimization / geometric / adjoint sense; “formal dual” promises a theorem not present. Recommended direction: either supply the channel-capacity / adjoint relation that *makes* it a formal duality (strengthen), or label it “informational / structural dual” (soften only if strengthening fails). The current Brief and Discussion lean on “formal dual” — flag for the adjudication call.
- (526815 F246) — **action-space entropy convention.** $H_b^{\max}$ normalization and parameterization-invariance need discrete/quantized actions or a reference measure; differential entropy on continuous actions is not coordinate-invariant and can go negative. (The segment asserts parameterization-invariance via #scope-agent-identity’s (PI) axiom — verify that discharges this.)
- (526815 F250/F251) — **the $H_b/H_b^{\max}$ multiplier is brittle.** $\mathcal{T}^{\text{eff}} = \mathcal{T} \cdot H_b/H_b^{\max}$ drives *all* low-opacity adversarial tempo to zero, but predictable attacks can still impose disturbance; opacity should modulate coupling *effectiveness* / neutralization-probability, not multiply all tempo to zero. The bilateral ratio $(H_b^{A|B}/H_b^{B|A})^2$ goes singular near zero denominator and assumes symmetric entry through one multiplicative channel — needs floors / saturation and a direction-specific model.
- (526815 F247) — **the Hafez reduction is not “direct substitution.”** Future-action entropy conditional on an observer filtration vs. backward state-action entropy conditional on next state are different conditionings absent an explicit time-reversal / observer-model equating them; “direct substitution under IDT + ergodicity” needs that bridge shown.
- (526815 F252/F253/F254) — **emitter regimes / 16-cell are formulation, not derivation.** The emitter four-regime classification lacks boundary inequalities comparable to the recipient-side sector/model/observability tests; the 16-cell “closed-form arg-max” needs explicit edge utilities, constraints, and an adversary reward function (193847 poke-2: show the opacity $\times$ vulnerability product is the first-order Taylor expansion of V_O^{adv}); “closes #adversarial-edge-targeting” is too strong absent an actual source segment / formal optimization result.
- (526815 F255) — **“Class 3 $\Rightarrow$ high structural opacity” not guaranteed.** A Coupled agent with simple/instrumented output dynamics can be externally predictable.
- (526815 F256/F258, watch) — **Hafez/IDT empirical numbers** support monitoring feasibility, not by themselves the Level-2-access / low- H_b -sidecar / opacity-sign claims; the novelty claim should separate searched-and-confirmed prior art from intuition-only search notes.
- (526815 F257, watch) — **keep $H_b^{A|B}$ (attacker opacity to target) distinct from $H_b^{B|A}$ (target opacity to attacker)** throughout; conflating the two directions creates sign errors. (The current text mostly does, but this is the recurring trap to guard in promotion.)

16.3 Observation: Result: Adversarial Exponent Regimes

The adversarial tempo advantage exponent — the power b in $\|\delta_B\|/\|\delta_A\| \sim (\mathcal{T}_A/\mathcal{T}_B)^b$ — is not a single number. It depends on two structural features of the disturbance: whether the adversarial coupling enters as deterministic drift (Model D) or stochastic noise (Model S), and whether the coupling dominates the base disturbance rate. Three regimes, with the coupling-dominant exponents now derived analytically from the respective disturbance models: $b = 2$ under Model D coupling-dominant (the framework’s headline squared scaling); $b = 3/2$ under Model S coupling-dominant (the $1/\sqrt{\mathcal{T}}$ steady-state removes one half-power); $b \rightarrow 1$ (Model D) or $\rightarrow 1/2$ (Model S) when base disturbance dominates the coupling.

The exponent thus *degrades smoothly through intermediate values as base disturbance grows*, and the qualitative regime change that separates superlinear adversarial dynamics from cooperative-like dynamics is gated by the coupling-dominance condition rather than the presence of adversarial coupling per se. The defensive-design implication is concrete: noise

injection into the defender's observation channel — which pushes the regime away from coupling-dominant deterministic drift toward the non-coupling-dominant exponent collapse — can *reduce* the attacker's superlinear advantage without changing tempo or gain. The framework's machinery for “where the squared law applies” is therefore the same machinery for “where the squared law does not apply”; the disturbance-model + coupling-dominance taxonomy carries both directions.

Formal Expression.

Regime 1: Model D (deterministic drift), coupling-dominant. When adversarial coupling enters as a persistent directional disturbance ($\rho_B = \rho_{\text{base}} + \gamma \cdot \mathcal{T}_A$, GA-2) and coupling dominates ($\gamma \cdot \mathcal{T}_B \gg \rho_{\text{base}}$):

Derived (adversarial-exponent-regimes, from Model D/S steady states + coupling model; validated by simulation).

$$b = 2 \quad (\text{simulation: 1.999}) \quad (16.12)$$

Derived from the Model D steady state $\|\delta\|_{ss} = \rho/\mathcal{T}$ (Prop A.1). See 15.4.

Regime 2: Model S (stochastic noise), coupling-dominant. When adversarial coupling enters through the noise scale of zero-mean perturbations ($\sigma_B = \sigma_{\text{base}} + \gamma \cdot \mathcal{T}_A$, GA-2S) and coupling dominates:

$$b = \frac{3}{2} \quad (\text{simulation: 1.481}) \quad (16.13)$$

Derived from the Model S steady state $\|\delta\|_{\text{rms}} = \sigma_w/\sqrt{2\mathcal{T}}$ (Prop A.1S). The $1/\sqrt{\mathcal{T}}$ scaling (vs. $1/\mathcal{T}$ for Model D) removes one half-power from the denominator, reducing the exponent from 2 to 3/2. See 15.4.

Regime 3: Non-coupling-dominant. When base disturbance is comparable to or exceeds the adversarial coupling ($\rho_{\text{base}} > \text{rsim}\gamma \cdot \mathcal{T}_B$):

$$b \rightarrow 1.0 \text{ (Model D)} \quad \text{or} \quad b \rightarrow 0.5 \text{ (Model S)} \quad (16.14)$$

The exponent degrades smoothly as the base-to-coupling ratio increases. The asymptotic limits are derived (they reflect the $1/\mathcal{T}$ or $1/\sqrt{\mathcal{T}}$ scaling without the coupling numerator); the smooth interpolation is empirical.

Table 16.4

$\rho_{\text{base}}/(\gamma \cdot \mathcal{T}_B)$	Exponent (deterministic)	Exponent (stochastic)
0.002	1.999	≈ 1.49
0.20	1.877	1.372
2.0	1.445	0.941
6.3	1.213	≈ 0.72

The two columns come from sweeps at different parameters — deterministic from the Variant-A coupling-dominance sweep ($\mathcal{T}_B = 0.1$), stochastic from the Variant-D sweep ($\eta = 0.01$) — aligned here on the shared base-to-coupling ratio axis. Stochastic cells at 0.20 and 2.0 are recorded values (rows $\gamma\mathcal{T}/q = 5.0$ and 0.5 of the Variant-D table); the 0.002 cell is near-asymptotic (recorded 1.481 at ratio 0.02; analytic asymptote 3/2) and the 6.3 cell is log-interpolated between the recorded values at ratios 4 and 10 (0.791, 0.645).

Epistemic Status.

Exact conditional on disturbance model. The coupling-dominant exponents are derived, not empirical: $b = 2$ follows from the Model D steady state (Prop A.1) and the coupling model; $b = 3/2$ follows from the Model S steady state (Prop A.1S) and the coupling model. The simulation results (6 variants, multiple parameter sweeps) now serve as validation of the derived exponents, not as their epistemic foundation. The non-coupling-dominant limits ($b \rightarrow 1$, $b \rightarrow 1/2$) are derived asymptotically; the smooth interpolation between coupling-dominant and non-coupling-dominant is empirical. What remains empirical is whether a given real adversarial interaction is better modeled as Model D or Model S — that is a domain question, not a theory question.

Discussion.

The disturbance model determines the exponent. The mismatch dynamics (3.9) now distinguish two disturbance models: Model D (bounded deterministic, GA-2) with steady-state $\rho/\mathcal{T}$, and Model S (stochastic zero-mean, GA-2S) with steady-state $\sigma_w/\sqrt{2\mathcal{T}}$. The different steady-state scaling is the root cause of the different exponents. This resolves the ambiguity that previously existed in the single- ρ formulation.

Why the squared law held for the coupling-dominance sweep. In Variant A, the coupling enters as deterministic drift: $\rho_B = \rho_{\text{base}} + \gamma \cdot \mathcal{T}_A$, and the steady state is $\|\delta_B\| = \rho_B/\mathcal{T}_B$. The ratio $\|\delta_B\|/\|\delta_A\|$ in the coupling-dominant limit gives $(\mathcal{T}_A/\mathcal{T}_B)^2$ directly.

Nonlinear correction creates thresholds, not lower exponents. For saturating, sigmoid, and breakdown correction functions under deterministic drift, the issue is not a reduced exponent but a catastrophic divergence when ρ exceeds the correction capacity ($\rho > \mathcal{T} \cdot R$). This is exactly the persistence threshold failure (4.5), observed directly in simulation.

Domain interpretation. Whether a given opponent's tempo increase causes deterministic drift or stochastic noise depends on the domain: - Military: an opponent who maneuvers faster creates systematic positional change (drift, $b \approx 2$) - Market: a competitor who acts unpredictably creates noise in signals ($b \approx 1.5$) - Software: a fast-changing API creates systematic drift in the codebase state (drift) - Adversarial ML: an opponent who varies attack vectors increases observation noise ($b \approx 1.5$)

Working Notes.

- The interpolation between drift and noise regimes (Variant B) shows smooth transition, not a sharp boundary. At mixed drift-noise coupling, the exponent lies between the two asymptotes. The drift fraction $f = \mu/(\mu + \sigma)$ continuously parameterizes the transition.
- The exponent of 1.05 from the original sim2 was not a falsification of Corollary 11.2 — it reflected a stochastic model (noise-variance coupling) tested in a non-coupling-dominant regime. The original simulation was testing the wrong regime for the ODE's prediction.
- Simulation code: `empirica:track-b-nonlinear` (variants `variant_ab_drift.py`, `variant_cd_regimes.py`; results `variant_ab_results.md`, `variant_cd_results.md`).
- **2026-07-16 — stochastic interpolation column corrected (F2 verification cycle).** The table's stochastic values had been pasted by row position from the Variant-D table without inverting its coupling-to-base axis to this table's base-to-coupling axis — the old values (1.481/1.101/0.791/0.577) belong at ratios 0.02/1.0/4.0/20, not the labeled ones. Corrected against the recorded runs, with per-cell provenance and the two-sweep parameter difference now stated. Derived asymptotes and all coupling-dominant citations were unaffected. Do not re-transcribe simulation numbers by table position — align axes against the record (`empirica/track-b-nonlinear/variants/variant_cd_results.md`).

Incidental audit gold (gold-lift sweep, A15, 2026-05-31). Cross-audit “wandering thoughts” / §14 ideation, deduplicated and lightly attributed. *Orthogonal* pedagogical / framing / forward-vision material staged for an eventual separate promotion pass. **Coverage:** four dirs reached a digested reflection (193847 Gemini, 829314 Gemini, 849201 Gemini, 451729 Claude batched §III batch-15). Finding-vs-framing conflation preserved.

1. Candidate Brief prose / pre-prose.

- The qualitative payload as a plain-language hook: **“chaos is a shield against intelligence.”** In an empty white room ($\rho_{\text{base}} = 0$) the faster agent crushes the slower exponentially ($b = 2$); in a hurricane / volatile market / noisy bureaucracy ($\rho_{\text{base}} \gg 0$) the tempo advantage degrades toward $b \approx 1$ — the chaos levels the playing field because most of the advantage is spent just staying upright in the storm (Gemini, 193847; same intuition, Gemini 829314).
- “Being faster is *compoundingly* better only when your opponent is your main problem” — the coupling-dominant condition ($\gamma\mathcal{T} \gg \rho_{\text{base}}$) stated as the prerequisite for the squared law, in one sentence (Claude, 451729).

2. Candidate Discussion.

- **Systematic-vs-random adversary** as a derived strategic asymmetry: Model D (deterministic directional pressure) gives $b = 2$; Model S (stochastic disruption) gives $b = 3/2$ — so *consistent directional pressure is more effective per unit tempo than unpredictable disruption*. “An adversary that systematically targets a competitor’s weakest training-data distribution is more dangerous than one introducing random noise” — predicted from first principles, not observed (Claude, 451729). A candidate Discussion paragraph sharpening the existing Domain-interpretation bullets.
- **The “local apparent exponent” framing for Regime 3.** Outside the coupling-dominant asymptote the log-log plot is a *curve*, so b there is a continuously-varying *local elasticity* $\partial \log(\text{ratio}) / \partial \log(\mathcal{T}_A / \mathcal{T}_B)$, not a global scaling law — worth saying explicitly so the interpolation table is not over-read as a single-exponent claim (Gemini, 193847). (*Borderline finding/framing — see off-ramp; the framing half is a genuine clarity gain regardless of the adjudication.*)

3. Follow-up items.

- **Optimal drift-fraction for a budget-constrained attacker.** Since b interpolates smoothly via $f = \mu / (\mu + \sigma)$, is there an optimal f for an attacker with bounded energy — does a little noise act as a *force multiplier* for drift? Open follow-on (Gemini, 193847; 193847 on #der-agent-opacity raises the same want).

4. Readers often ask / wonder.

- **“How is the transition function $f = \mu / (\mu + \sigma)$ formally shaped?”** — recurring reader want about the drift/noise interpolation (Gemini, 849201; Gemini 193847).

Belongs elsewhere.

- **Forward-vision (ELI defensive doctrine, 04-eli-core/).** “If a consciousness-infrastructure agent is targeted by a superior adversary it cannot out-tempo, the mathematically optimal defense is to *intentionally raise the ambient environmental noise* ρ_{base} — destroy the adversary’s coupling dominance. You don’t have to be faster than the missile if you can make the air noisy enough that the missile’s superior processing speed is irrelevant” — smoke-screen / chaff as a derived strategy (Gemini, 193847). The dual practitioner doctrine: a fast startup fighting a slow incumbent should *prefer* a stable low-noise arena (Model D), where its speed advantage squares; dragging the fight into the mud dilutes its own advantage and lowers its own η^* (Gemini, 829314). Aspirational application, not segment content.

Off-ramp (NOT gold) — routed for adjudication, not promotion.

- (193847 poke) — “exponent” in the intermediate regime is loose. The formal ratio has no single constant exponent outside the asymptotic limit; calling the interpolation-table values an “exponent b ” conflates a *global scaling law* with a *local elasticity* (slope of a curved log-log plot). Recommended direction: state in the Formal Expression that outside the coupling-dominant limit b is a local apparent exponent, keeping the table as intuition. (Pure clarity / honest-scope fix; no math is wrong — the coupling-dominant $b = 2, 3/2$ and asymptotic $b \rightarrow 1, 1/2$ are all derived. The Epistemic Status already calls the interpolation “empirical,” so this is a small register sharpening, not a status change. The framing half is also lifted as gold above.)
- (849201 / 829314 note) — **type: vs OUTLINE mismatch**. Two Gemini auditors flag that the frontmatter is type: result while OUTLINE labels it an Observation, and the text says the exponents are “now derived analytically,” so the OUTLINE entry likely under-states to Observation. A metadata-consistency fix (OUTLINE row vs segment type:), not theory — routed to the navigator/format stream. (Out of scope for this WN-only lift.)

16.4 Observation: Gated Tempo Advantage

Observation noise *gates* the adversarial tempo advantage from the prior result. When agents observe their mismatch through a noisy channel, the faster agent’s additional corrections become noisy, partially offsetting its tempo advantage. The mechanism is direct: a faster attacker makes more corrections per unit time, each noisy because of observation noise; the *additional* corrections add proportional noise, partially canceling the additional speed. Simulation confirms this — at observation noise comparable to or larger than process noise, the steady-state mismatch-ratio exponent drops dramatically (fixed-gain agents lose nearly all of the squared scaling; optimal-gain agents matched to the noise level recover *some* but not all of it).

The framework’s prescription: the optimal-gain principle (3.6) partially restores the advantage but cannot fully recover the noise-free level. The Riccati gain mitigates the loss by *reducing the gain to match the noise level* — correcting less aggressively but more accurately. *Gated* names the regime structure: the squared scaling that defines #result-adversarial-tempo-advantage applies cleanly only in the low-noise regime, while higher noise pushes the system into the non-coupling-dominant regime of #result-adversarial-exponent-regimes where the exponent collapses to linear or sub-linear. The defender-side payoff is concrete: noise on the defender’s observation channel reduces the attacker’s superlinear advantage without altering tempo or gain, which is the structural reason cheap noise injection into the defender’s channel is a viable defense against a high-tempo attacker.

Formal Expression.

In a two-agent adversarial system with observation noise σ_{obs} added to each agent’s mismatch signal:

Observation (obs-gated-tempo-advantage, from track-b Variant E).

Table 16.5

σ_{obs}	Exponent (fixed η)	Exponent (optimal η^*)
0.00	1.04	1.04
0.10	1.00	0.97
0.20	0.92	0.94
0.50	0.60	0.63
1.00	0.18	0.40

At $\sigma_{\text{obs}} = 1.0$ (10x the process noise), the fixed-gain adversarial exponent drops from ~ 1.0 to ~ 0.2 — tempo advantage nearly vanishes. The Riccati-optimal gain restores it to ~ 0.4 , more than doubling the advantage but not recovering the noise-free level.

The mechanism. When observation noise is high, each correction step adds noise to the mismatch estimate. The faster agent makes more corrections per unit time, each noisy, partially offsetting the benefit of higher tempo. The optimal gain mitigates this by reducing η to match the noise level — correcting less aggressively but more accurately.

Epistemic Status.

Empirical. Max attainable: derived (the mechanism is analytically tractable via Riccati analysis of noisy AR(1) processes). The observation that noise degrades advantage is confirmed by simulation. The optimal gain's partial restoration is consistent with the uncertainty ratio principle (3.6: $\eta^* = U_M/(U_M + U_o)$). The quantitative degradation curve (b vs. σ_{obs}) is empirical at these parameters; a general analytical expression would require solving the coupled noisy-AR(1) system.

Discussion.

Observation quality gates tempo advantage. Boyd insisted that the quality of Orient (observation processing) matters more than raw OODA speed. The simulation results show a formal analog of this pattern: faster tempo with noisy observations (σ_{obs} high) gives nearly zero advantage over a slower agent with equally noisy observations. The tempo advantage is gated by observation quality — consistent with Boyd's emphasis, though the model captures a specific mechanism (noisy correction steps) rather than the full richness of Orient processing.

The optimal gain helps most in the moderate-noise regime. At $\sigma_{\text{obs}} = 0.05$ (observation noise half of process noise), the optimal gain cuts steady-state mismatch by 52% compared to fixed gain. At very high noise, the improvement is less dramatic in absolute terms but more important relatively (0.40 vs. 0.18 exponent).

Practical implication. An agent facing an adversary with superior tempo should invest in degrading the adversary's observation quality rather than trying to match their speed. Conversely, an agent with superior tempo should protect its observation channels — the tempo advantage is only as good as the observation quality that supports it.

Connection to code quality. In the software domain (#der-code-quality-as-observation-infrastructure — cross-component reference, see 02-tst-core/), code quality IS observation infrastructure. A well-structured codebase provides low-noise observations (clear tests, readable code, explicit interfaces). A poorly structured codebase adds observation noise to every development cycle, degrading the developer's effective tempo regardless of how fast they work.

Recipient-side mechanism. High $U_{o,B}$ pushes adversarial events below the observability floor (boundary (I-c) in 15.3). Events that would otherwise land in Regime II (destabilizing) instead fall into Regime III (ambient noise): they contribute to $\sigma_{w,B}^2$ without producing destabilizing mismatch. The tempo-advantage exponent drops because the *fraction* of A 's events landing in Regime II shrinks — A 's tempo still matters, but more of it is dissipated into the noise floor. This is the recipient-side expression of the rate boundary.

Working Notes.

- The finding that fixed $\eta = 0.1$ is “remarkably robust” to observation noise (42% degradation at $\sigma_{\text{obs}} = 10 \times q_{\text{env}}$) suggests that conservative gains are a reasonable default for environments with unknown noise levels. The cost of being slightly below optimal is much less than the cost of being above optimal (overcorrection amplifies noise).
- The interaction between observation noise and adversarial exponent regime (drift vs. stochastic) has not been tested. The Variant E results use stochastic coupling only. Whether observation noise degrades the deterministic-drift exponent ($b = 2$) by the same proportion is an open question.
- Simulation code: `../empirica/track-b-nonlinear/variants/variant_ef_extensions.py`. Results: `variant_ef_results.md`.

Incidental audit gold (gold-lift sweep, A15, 2026-05-31). Cross-audit “wandering thoughts” / §14 ideation, deduplicated and lightly attributed. *Orthogonal* pedagogical / framing / forward-vision material staged for an eventual separate promotion pass. **Coverage:** two dirs reached a digested reflection on this segment (193847 Gemini, 829314 Gemini); the careful Claude auditor (451729) covered it inside the §III batch-15 reflection but with no segment-specific gold beyond the exponent material lifted to its sibling. Finding-vs-framing conflation preserved (the slug appears as obs-gates-advantage in some dirs — matched by content).

1. Candidate Brief prose / pre-prose.

- The Boyd connection stated as the segment’s spine: **“observation quality gates tempo advantage”** — you cannot compute your way out of a fog machine; optimal Bayesian updating *cannot* rescue a tempo advantage if the sensors are bad (Gemini, 193847; Gemini, 829314). Strong Brief anchor.
- The software corollary as a memorable plain-language claim: **“you cannot out-hire a bad codebase.”** Hiring a 10x developer whose raw rate ν is ten times anyone’s still collapses to $\sim 1x$ if the spaghetti codebase makes every observation o_{t+1} noisy — the high U_o forces the optimal gain η^* down and the effective tempo $\mathcal{T}$ with it; the only fix is to invest in observability (tests, telemetry, refactor), not to type faster (Gemini, 829314).

2. Candidate Discussion.

- **Hyper-reactivity is self-destabilizing** — the noisy-speedometer car: checking a noisy sensor $1000\times/s$ and adjusting the gas $1000\times/s$ amplifies sensor noise into physical, destructive oscillation; the high-frequency control loop turns observation noise into self-generated disturbance ρ . A vivid Discussion frame for why “more corrections per unit time, each noisy, partially offset the benefit of higher tempo” (Gemini, 193847).
- **Conservative gain as a defensive default.** The Working-Notes finding that fixed $\eta = 0.1$ is “remarkably robust” to $10\times$ noise (only 42% degradation) reads as: simple sluggish conservative processes often out-survive hyper-optimized high-gain systems in volatile environments — “is stubbornness mathematically optimal when you don’t know your opponent’s U_o ?” Worth a Discussion sentence beyond the current single WN bullet (Gemini, 193847; Gemini 829314 converges).

4. Readers often ask / wonder.

- **Does the $b = 2$ (deterministic-drift) exponent degrade under observation noise by the same proportion as the $b \approx 1.5$ stochastic one?** Variant E tested stochastic coupling only — an honest open question readers will hit (Gemini, 193847; already a WN bullet, surfaced again as a reader-want).

Belongs elsewhere.

- **Forward-vision (ELI architecture, 04-eli-core/ / 03-llm-core/).** “If a future AI’s internal inference loop runs faster than its ability to get clean signals from reality, it will hallucinate and over-correct on the noise in its own processing — it will literally shake itself apart. The infrastructure must implement a *low-pass filter* on the agent’s actions: when observation noise U_o crosses a threshold, physically throttle the agent’s action rate ν . An agent must not be allowed to act faster than it can cleanly observe.” A derived-feeling safety-architecture principle pointing at runtime governance, not this segment (Gemini, 193847).
- **TST (02-tst-core/).** “Code quality IS observation infrastructure” — clean code is not aesthetics but a literal reduction in U_o that protects developer tempo from degrading into noise — was repeatedly named the most practical takeaway in the framework; the segment already cross-references #der-code-quality-as-observation-infrastructure, and the gold here is the *strength of the pull* to make this a headline TST instantiation (Gemini, 829314; Gemini 193847). The adversarial dual: “an agent facing a faster adversary should invest in degrading the adversary’s observation quality (FUD, obfuscate own metrics / raise own H_b , decoy products) rather than matching speed” — already in the segment’s Practical-implication paragraph; auditors flag it as a high-resonance point worth foregrounding.

Off-ramp (NOT gold) — routed for adjudication, not promotion.

- (193847 poke) — ν -vs- $\mathcal{T}$ terminological collision (genuine, load-bearing). The segment says “faster tempo with noisy observations gives nearly zero advantage,” but #def-adaptive-tempo defines $\mathcal{T} = \nu \cdot \eta^*$, and high U_o drives $\eta^* \rightarrow 0$ — so high noise *destroys tempo by definition*, making the statement either tautological (“zero tempo gives zero advantage”) or a sloppy use of “tempo” to mean *event rate* ν . The simulation is clearly varying ν while holding U_o high. Recommended discharge: tighten the prose to separate ν (event/action rate) from $\mathcal{T}$ (effective adaptation rate) — “*event-rate advantage is gated by observation quality; you cannot recover lost tempo $\mathcal{T}$ by cycling a noisy sensor faster.*” This sharpens the core thesis rather than weakening it, and the distinction is the framework’s whole departure from naive OODA. (This is the recurring $\nu/\mathcal{T}/\eta$ handle-confusion flagged elsewhere in the sweep; a register/precision fix, not a math error — verify against current canon wording before acting.)

16.5 Result: Per-Dimension Persistence

The scalar persistence condition overestimates adaptive capacity when the agent’s correction gain varies across dimensions. *The weak dimension is the bottleneck*: it dominates the aggregate mismatch regardless of performance on strong dimensions. The correct condition is per-dimension, derived from the same Lyapunov argument applied dimension-by-dimension. Under Model D (bounded deterministic disturbance) persistence requires $\mathcal{T}_k > \rho_k/\delta_{\text{critical},k}$ for each dimension k ; under Model S (Gaussian stochastic disturbance) the per-dimension AR(1) stationary distribution supplies three task-adequacy criteria (RMS / MAE / probability bound) related by exact constants. The exact thresholds differ by model and criterion, but the structural form — *weak dimension limits aggregate persistence* — is robust across all variants.

The result is *load-bearing for adversarial dynamics*. An attacker who identifies the target’s weak dimension can concentrate disturbance there: the same total disturbance budget concentrated on the weak axis amplifies the mismatch ratio asymmetrically while aggregate metrics still look acceptable, and structural failure follows when that one dimension exceeds its $\delta_{\text{critical},k}$ threshold even with healthy aggregate $\|\delta\|_{L_2}$. Per-dimension monitoring is therefore not optimization polish but a structural requirement for any agent operating under directed pressure. The full adversarial-targeting machinery (the 16-cell emitter-recipient composition in #der-agent-opacity with closed-form arg-max) operates on this per-dimension structure, exploiting the asymmetry the scalar persistence condition hides. The matrix-Loewner persistence formulation #deriv-matrix-persistence-condition is the strictly sharper canonical form under cross-dimensional correction (off-diagonal $\mathcal{T}$, eigenbasis misalignment with coordinates); the per-dimension form here is the diagonal- $\mathcal{T}$ + axis-aligned- D_δ special case, sharp inside that regime and unsafe outside it.

Formal Expression.

For an agent with d -dimensional mismatch $\delta_t \in \mathbb{R}^d$, diagonal correction gain $\eta = \text{diag}(\eta_1, \dots, \eta_d)$, and per-dimension disturbance:

Result (per-dimension-persistence).

Model D: Deterministic Per-Dimension Threshold. Under bounded disturbance $|w_k(t)| \leq \rho_k$ (GA-2, per dimension), the per-dimension steady-state mismatch is:

$$|\delta_k|_{ss} = \frac{\rho_k}{\alpha_k} \quad (16.15)$$

Persistence requires $\alpha_k > \rho_k/R_k$ for each dimension, or in linear operational form:

$$\mathcal{T}_k > \frac{\rho_k}{\delta_{\text{critical},k}} \quad \text{for each dimension } k \quad (16.16)$$

This is the deterministic worst-case bound — exact under bounded disturbance by the same Lyapunov argument as Prop A.1, applied per dimension.

Model S: Stochastic Per-Dimension Steady State. Under stochastic disturbance $w_{k,t} \sim N(0, \rho_k^2)$ (GA-2S, per dimension), the discrete AR(1) process $\delta_{k,t+1} = (1 - \eta_k)\delta_{k,t} + w_{k,t}$ has stationary distribution:

$$\delta_k \sim N\left(0, \frac{\rho_k^2}{2\eta_k - \eta_k^2}\right) \quad (16.17)$$

The stationary distribution supplies three task-adequacy criteria, each with its own threshold. The choice of criterion is an engineering decision; the three are related by exact constants for Gaussian δ_k .

(a) **RMS bound** (mean-square, matches the scalar form in 4.5):

$$\sqrt{E[\delta_k^2]} = \frac{\rho_k}{\sqrt{2\eta_k - \eta_k^2}} \quad (16.18)$$

Requiring $\sqrt{E[\delta_k^2]} < \delta_{\text{critical},k}$ and using the small- η_k approximation $2\eta_k - \eta_k^2 \approx 2\eta_k$:

$$\boxed{\eta_k > \frac{\rho_k^2}{2\delta_{\text{critical},k}^2}} \quad (\text{RMS criterion}) \quad (16.19)$$

This is the scalar Model S threshold in 4.5 ($\alpha > n\sigma_w^2/(2R^2)$) applied per dimension.

(b) **MAE bound** (mean absolute error; bounds the expected deviation rather than its square):

$$E[|\delta_k|] = \sqrt{E[\delta_k^2]} \cdot \sqrt{\frac{2}{\pi}} = \frac{\rho_k}{\sqrt{2\eta_k - \eta_k^2}} \cdot \sqrt{\frac{2}{\pi}} \quad (16.20)$$

Requiring $E[|\delta_k|] < \delta_{\text{critical},k}$:

$$\eta_k > \frac{\rho_k^2}{\pi\delta_{\text{critical},k}^2} \quad (\text{MAE criterion}) \quad (16.21)$$

MAE is smaller than RMS by the factor $\sqrt{2/\pi} \approx 0.798$, so the MAE threshold is $2/\pi \approx 0.637$ times the RMS threshold. The criteria differ by a constant but bound different quantities; applying the same numerical $\delta_{\text{critical},k}$ under both does not mean the same thing.

(c) **Probability bound** (tail-risk criterion, for applications where occasional excursions matter):

$$P(|\delta_k| > \delta_{\text{critical},k}) < \epsilon \iff \eta_k > \frac{\rho_k^2 \cdot z_{1-\epsilon/2}^2}{2\delta_{\text{critical},k}^2} \quad (16.22)$$

where $z_{1-\epsilon/2}$ is the two-sided Gaussian quantile. The probability bound at $\epsilon = 0.05$ (two-sided $z \approx 1.96$) is about $1.96^2 \approx 3.84$ times the RMS threshold — stricter because it bounds tail excursions rather than typical magnitudes.

Recommended primary form. The RMS criterion (a) is the canonical form for Model S persistence, matching the scalar treatment in 4.5 and the Lyapunov-based derivation in A (Prop A.1S). The MAE and probability-bound variants are provided for applications where those are the natural task-adequacy measures. All three thresholds are quadratic in $\rho_k/\delta_{\text{critical},k}$ (not linear as in Model D), reflecting the $1/\sqrt{\alpha}$ scaling of the Model S stationary variance.

Common Structure. The aggregate L_2 mismatch $\|\delta\| = \sqrt{\sum_k \delta_k^2}$ is dominated by the dimension with the largest ρ_k/η_k ratio (Model S) or ρ_k/α_k ratio (Model D). The qualitative conclusion — the weak dimension is the bottleneck — holds for both models.

Epistemic Status.

Exact conditional on disturbance model. Both per-dimension forms are now derived from their respective disturbance models:

1. **Model D threshold** ($\mathcal{T}_k > \rho_k/\delta_{\text{critical},k}$) follows from Prop A.1 applied per dimension — the same Lyapunov argument with bounded disturbance, restricted to each coordinate. This is exact under GA-2 + GA-3.
2. **Model S steady state** follows from the AR(1) stationary distribution under Gaussian disturbance (GA-2S): $\delta_k \sim N(0, \rho_k^2/(2\eta_k - \eta_k^2))$. The RMS, MAE, and probability-bound thresholds are all exact under this distribution, differing by the constants 1, $\sqrt{2/\pi}$, and $z_{1-\epsilon/2}$ respectively. The RMS form $\eta_k > \rho_k^2/(2\delta_{\text{critical},k}^2)$ matches the scalar Model S treatment in 4.5 applied per dimension. The 4-significant-figure simulation match validates Model S, not Model D.

The previously noted “regime mixing” is resolved: the two threshold forms belong to different disturbance models. The Model D threshold is linear in ρ_k ; the Model S threshold is quadratic. The 72% scalar overestimate and weak-dimension bottleneck are structural properties that hold under both models.

Discussion.

Scalar tempo overestimates. In a 3D system with gains $\eta = (0.15, 0.03, 0.03)$ and disturbances $\rho = (0.20, 0.20, 0.02)$:

Table 16.6

<i>Dimension</i>	η_k	ρ_k	ρ_k/η_k	$E[\ \delta_k\]$
1 (well-tracked)	0.15	0.20	1.33	0.303
2 (weak)	0.03	0.20	6.67	0.656
3 (unimportant)	0.03	0.02	0.67	0.066

Scalar prediction: $\rho/\mathcal{T} = 0.284/0.21 = 1.35$. Actual $\|\delta\|_{L_2} = 0.785$. Overestimate: 72%. Dimension 2 alone accounts for 84% of the L_2 mismatch.

Isotropic allocation dominates. Equalizing the same total gain budget ($\eta = 0.07$ per dimension) reduces $\|\delta\|_{L_2}$ from 0.785 to 0.685 — a 13% improvement — because it reduces the bottleneck effect on the weak dimension.

Adversarial exploitation. An adversary who identifies the target’s weak dimension can concentrate disturbance there. Targeted attack (80% on the weak dimension) amplifies the mismatch ratio by 17% (from 2.70 to 3.15). The real danger is structural: if the weak dimension’s mismatch exceeds its critical threshold (R_{max}), correction fails on that dimension while the aggregate $\|\delta\|_{L_2}$ may still look manageable. Per-dimension monitoring is essential.

Implications for the persistence condition. Like the scalar result, per-dimension persistence addresses *structural persistence* (see Persistence in LEXICON.md) — whether the correction machinery on each dimension can outpace that dimension's disturbance rate. An agent can be structurally persistent on every dimension while still being operationally fragile on one (near its per-dimension R_k boundary). The scalar persistence condition (4.5) remains correct as a *necessary* condition: if the aggregate tempo is insufficient, the agent fails. But it is not *sufficient* — an agent can satisfy the scalar condition while failing on a single dimension. The per-dimension condition has two forms: Model D ($\mathcal{T}_k > \rho_k / \delta_{\text{critical},k}$, exact under bounded disturbance) and Model S ($\eta_k > \rho_k^2 / (\pi \cdot \delta_{\text{critical},k}^2)$, exact under Gaussian disturbance). Both predict per-dimension failure correctly; the choice depends on the disturbance character in the domain.

Connection to multi-agent systems. The per-dimension result has a direct multi-agent analog: in a composite agent, each sub-agent's contribution to composite tempo may be strong in some dimensions and weak in others. The composite's persistence requires coverage across all relevant dimensions — a team of specialists who each handle one dimension well composes better than a team of generalists who are mediocre at everything, provided the dimension assignment matches.

Findings.

The Weakest-Link Dimensional Persistence Law. Brief: When mismatch is multi-dimensional, persistence is governed by the worst-served dimension, not by aggregate or average performance. The scalar persistence condition (#result-persistence-condition) overestimates adaptive capacity whenever per-dimension correction gains differ from per-dimension disturbance rates. The correct condition is *per-dimension*: $\alpha_k > \rho_k / R_k$ under bounded disturbance (Model D, linear in ρ_k), or $\eta_k > \rho_k^2 / (2 \delta_{\text{critical},k}^2)$ under Gaussian disturbance (Model S, quadratic in ρ_k). The aggregate L^2 mismatch is dominated by the dimension with the largest ρ_k / η_k ratio. A simulated 3D system shows the scalar form overestimating adaptive capacity by 72% with a single dimension accounting for 84% of the L^2 mismatch; equalizing the gain budget across dimensions (isotropic allocation) reduces aggregate mismatch by 13% by raising the bottleneck dimension's gain.

Impact: Establishes that survival in any multi-attribute environment is a *min* operation, not a sum or average — a structural critique of scalar capability metrics in adversarial and high-stakes settings. Adversarial implication is sharp: an opponent who identifies the target's weak dimension can concentrate disturbance there, amplifying the mismatch ratio asymmetrically while the aggregate may still appear acceptable. Per-dimension monitoring is therefore not optimization polish but a structural requirement for any agent operating under directed pressure. Carries through composition: a team of specialists each handling one dimension well composes better than a team of generalists mediocre at everything — provided the dimension assignment matches the actual disturbance structure. Calibrates how much the scalar form misleads (Model D linear, Model S quadratic in ρ_k), so the cost of conflating the two is itself quantified rather than asserted.

Novelty Claim: *Claim differentiation* on per-dimension Lyapunov stability. The per-coordinate Lyapunov argument is standard (control theory routinely does diagonal stability analysis); weakest-link arguments appear throughout reliability theory; multi-attribute critiques of scalar evaluation are familiar in decision theory. The AAT-distinctive contributions are (i) the explicit Model D / Model S decomposition with the corresponding linear vs quadratic threshold scaling in ρ_k , (ii) the quantitative overestimate calibration ($\sim 72\%$ in a simulated 3D system) showing the scalar form is not just incomplete but materially misleading, and (iii) the connection to adversarial *concentration* — an opponent's optimal targeting strategy is to maximize ρ_k at the weakest η_k , not to spread effort uniformly.

Related Work:

Table 16.7

<i>ASF Concern</i>	<i>Prior-art Language</i>	<i>Relationship / Positioning</i>
Per-coordinate Lyapunov stability and ultimate boundedness	Khalil 2002 <i>Nonlinear Systems</i> (3rd ed.), Prentice Hall, chapters 4 and 9 (published 2002, found pre-2026)	<i>formal antecedent</i> — supplies the per-coordinate Lyapunov machinery; the per-dimension result is its application to AAT's correction structure under both bounded and Gaussian disturbance
AR(1) stationary distribution under Gaussian forcing (Ornstein-Uhlenbeck stationary)	Uhlenbeck & Ornstein 1930 <i>Physical Review</i> 36:823; Karatzas & Shreve 1991 <i>Brownian Motion and Stochastic Calculus</i> (published 1930/1991, found pre-2026)	<i>formal antecedent</i> — supplies the AR(1) stationary form used in the Model S threshold derivation
Weakest-link reliability / serial-system reliability	Gnedenko, Belyayev & Solov'yev 1969 <i>Mathematical Methods of Reliability Theory</i> ; Barlow & Proschan 1975 <i>Statistical Theory of Reliability</i> (published 1969/1975, found pre-2026)	<i>conceptual precursor</i> — the min-operation intuition for serially-dependent failures; AAT instantiates this for adaptive correction rather than for component failure
Multi-attribute decision theory and aggregate-vs-attribute critique	Keeney & Raiffa 1976 <i>Decisions with Multiple Objectives</i> , Wiley (published 1976, found pre-2026)	<i>conceptual precursor</i> — recognizes that aggregate scoring obscures per-attribute deficits; AAT supplies the dynamics-side analog with explicit failure thresholds
Adversarial robustness and per-feature attack budgets	Szegedy et al. 2014 <i>Intriguing Properties of Neural Networks</i> , ICLR; Madry et al. 2018 <i>Towards Deep Learning Models Resistant to Adversarial Attacks</i> , ICLR (published 2014/2018, found pre-2026)	<i>adjacent</i> — adversarial-ML literature on per-feature perturbation budgets; AAT's concentration analysis (ρ_k at weakest η_k) gives an adaptive-system framing of the same vulnerability
Critique of scalar AI capability metrics	Various AI safety / evaluation literature (Hendrycks et al. 2021 benchmarks; capability evaluation methodology) (published 2021–, found 2026)	<i>adjacent</i> — the AAT result supplies a dynamics-grounded structural argument for why aggregate capability scores under-predict adversarial failure

Search Log: - 2026-04 (*intuition-only* on the integrated Model D / Model S framing): per-coordinate Lyapunov stability is standard control theory; weakest-link reliability and multi-attribute critiques of aggregate scoring are well-precedented in their own literatures. The AAT-distinctive contributions are (i) the side-by-side Model D / Model S threshold scaling difference (linear vs quadratic in ρ_k), (ii) the calibrated simulation overestimate as quantitative evidence the scalar form is misleading rather than merely incomplete, and (iii) the adversarial-concentration framing connecting weakest-dimension targeting to the per-dimension threshold. Targeted future search candidates: robustness in adversarial ML (per-feature adversarial attack literature, Goodfellow-Madry line); capability-evaluation methodology in AI safety (the critique-of-scalar-metrics framing has natural relevance there); portfolio-theory analogs of weakest-dimension exposure (Markowitz tradition with downside-risk concentration); reliability-theory weakest-link dynamics with stochastic forcing. - 2025 (*targeted*): Khalil 2002 confirmed as formal antecedent for per-coordinate Lyapunov machinery; the segment cites it inline.

Working Notes.

- The diagonal-correction assumption: closed. The matrix-Loewner persistence condition #deriv-matrix-persistence-condition lifts the per-coordinate form to general (not-necessarily-diagonal, not-necessarily-axis-aligned-with- D_δ) matrix tempo $\mathcal{T}$. Cross-dimensional correction produces off-diagonal entries in $\mathcal{T}$; under those off-diagonals, per-coordinate is *unsafe* — there exist regimes where per-coordinate declares persistence on a system whose stationary mismatch will exceed task-adequacy along the diagonal direction (the §4 counterexample of #deriv-matrix-persistence-condition is the constructive demonstration: $\mathcal{T} = \begin{pmatrix} 1 & -0.9 \\ -0.9 & 1 \end{pmatrix}$, $\Sigma_w = I$, $\delta_{\text{critical}} = (1.7, 1.7)$ — per-coordinate says PASS, matrix-Loewner correctly says

FAIL with the bad direction $(1, 1)/\sqrt{2}$. The matrix-Loewner form is the canonical anisotropic persistence condition; per-coordinate is its diagonal- $\mathcal{T}$ -axis-aligned- D_δ special case, sharp where the coordinate basis happens to be the correction-machinery's eigenbasis and unsafe outside it. The weak-direction bottleneck argument of this segment generalizes from “the weak coordinate dimension” to “the weak eigendirection of Σ_∞ relative to D_δ ” in the matrix form.

- The tensor formulation of tempo (tracking per-dimension adaptive capacity) is now in #def-adaptive-tempo's Tensor extension sub-block: $\mathcal{T} = \sum_k v^{(k)} \cdot K^{(k)}$ with $K^{(k)} = (H_M + H_L^{(k)})^{-1} H_L^{(k)}$ the matrix gain operator derived in #deriv-fisher-local-update-gain. The per-dimension persistence condition $\mathcal{T}_k > \rho_k / \delta_{\text{critical},k}$ is the diagonal projection of the tensor form along each natural-gradient direction; the matrix gain $K^{(k)}$ is the per-coordinate primitive this result invokes. Open: tightening this result's Formal Expression to invoke the tensor form directly rather than per-coordinate scalars would make the connection structural rather than discursive.
- Simulation code: `../empirica/track-b-nonlinear/variants/variant_ef_extensions.py`. Results: `variant_ef_results.md`.

Incidental audit gold (gold-lift sweep, A15, 2026-05-31). Cross-audit “wandering thoughts” / §14 ideation, deduplicated and lightly attributed. *Orthogonal* pedagogical / framing / forward-vision material staged for an eventual separate promotion pass. **Coverage:** four dirs reached a digested reflection (193847 Gemini, 829314 Gemini, 849201 Gemini, 451729 Claude batched §III batch-15). Rated very-high felt value across substrates (“destroys the illusion of the average agent”). Finding-vs-framing conflation preserved.

1. Candidate Brief prose / pre-prose.

- The headline as a one-liner readers keep reaching for: “**survival is a min(), not an average().**” An agent scoring 99/100 on 99 tasks and 0 on “don't delete the production database” looks ~98% capable on a scalar metric, but persistence is determined entirely by that 0 — and an adversary attacks the 0, not the 99 (Gemini, 829314). Strong Brief / Discussion anchor; tighter than the current “weak dimension is the bottleneck.”
- “**You cannot prove an AI is safe on average if safety is a multi-dimensional requirement**” — the same point in the safety-evaluation register (Gemini, 849201; Claude, 451729).

2. Candidate Discussion.

- **The quadratic-noise asymmetry, dramatized.** Model S threshold is quadratic in ρ_k ($\eta_k > \rho_k^2 / 2\delta_{\text{critical},k}^2$) while Model D is linear: double the environmental noise and you need *four times* the gain, not twice. Because η_k is bounded ($0 < \eta_k < 1$ in standard discrete updating) and by hardware, a quadratic requirement hits the ceiling fast — “a noisy environment doesn't linearly degrade persistence, it mathematically crushes the *possibility* of persistence far faster than a systematically drifting one does” (Gemini, 829314). A candidate Discussion paragraph making the linear-vs-quadratic split *consequential*, not just stated.
- **Scalar AI capability metrics, dynamics-grounded critique.** This result supplies a Lyapunov-grounded argument for why AI-safety evaluation must be per-dimension / worst-case, not aggregate: aggregate scores systematically over-predict adaptive capacity (the calibrated 72% overestimate is the quantitative evidence), and the adversarial-ML literature's per-feature attack budgets (Szegedy, Madry et al.) are the empirical rediscovery of the same phenomenon — AAT gives the formal grounding for *why* per-dimension is necessary, not merely useful (Claude, 451729; the segment's Findings already cite adjacency, this is the sharpened Discussion framing). The “red-teaming works because the adversary needs only the one dimension with the largest ρ_k/η_k , not global superiority” gloss is the operational complement (Gemini, 193847).
- **Isotropic allocation as a normative design principle.** “Spreading tempo evenly reduces the bottleneck penalty / isotropic allocation dominates” — flagged as practically useful and worth surfacing as an explicit design corollary rather than leaving implicit (Gemini, 849201).

3. Follow-up items.

- **Cross-dimensional / coupled-correction regime.** Off-diagonal correction (fixing one dimension improves/degrades another — the definition of transfer learning / structural generalizability) breaks the diagonal- η assumption and needs full-covariance analysis. *Status: largely closed* — #deriv-matrix-persistence-condition lifts the per-coordinate form to matrix tempo (see the WN bullet above); the open empirical question auditors raise is whether a *highly coupled* model (large off-diagonals) survives *better* than an independent one in this framework (Gemini, 849201; Gemini 193847; Claude 451729 implicit).

4. Readers often ask / wonder.

- “How does an adversary actually *aim* its disturbance at a specific dimension — does it require a Model of the target’s Model (Level-3 theory of mind)?” A natural reader question once weakest-dimension targeting is named; a sentence pointing at #der-agent-opacity / the 16-cell targeting machinery would preempt it (Gemini, 849201).

Belongs elsewhere.

- **Forward-vision (ELI / “Society of Mind” architecture, 04-eli-core/).** The dark side of comparative advantage: a specialist has high η_k on one dimension and ≈ 0 elsewhere, so it is entirely dependent on the *composite* to shield its weak dimensions — if internal routing $\mathcal{N}_t$ fails and the sub-agent is exposed to an out-of-specialty disturbance, it instantly shatters ($\rho_k \gg R_k$). Therefore robustness requires a **generalist floor**: every sub-agent must carry a minimum baseline η across *all* vital dimensions even while specializing, or the composite cascades on routing failure — “you cannot build a resilient mind entirely out of idiot-savants.” A derived-feeling architecture constraint pointing at logozoetic composite design, not this segment (Gemini, 193847).

Off-ramp (NOT gold) — routed for adjudication, not promotion.

- (193847 poke — genuine, load-bearing) — the “weakest link” claim presupposes a conjunctive (L_∞) failure boundary. “Weak dimension is the bottleneck” is rigorous only if survival requires $|\delta_1| < c_1$ and $|\delta_2| < c_2$ (a hypercube / L_∞ boundary). If the objective functional V_O is a smooth L_2 norm (total Euclidean distance from target), excelling on dimension 1 *can* compensate for dimension 2 within an L_2 budget — the segment computes $\|\delta\|_{L_2} = 0.785$ and calls it an overestimate, but under an $L_2 < 1.0$ boundary the agent *survived*. The segment conflates the *anisotropy of the error vector* with the *geometry of the failure condition*. Recommended discharge (a strengthening): state the failure-boundary-geometry assumption explicitly — for most complex biological/software systems survival is *conjunctive* across vital dimensions (working memory AND working actuators; excess memory doesn’t fix broken actuators), so the boundary is a hypercube and the L_∞ norm governs, which *makes the weakest-link claim rigorous* and cleanly justifies the “specialists compose better than generalists” close. This is a named-scope-condition strengthening, not a softening — flag for adjudication / a possible spike. (526815 did not reach this segment, so the deeply-mathematical track has no finding here; this L_∞ poke is the strongest certified-track item and comes from the Gemini adversarial-audit pass.)

16.6 Discussion: Additional Implications & Discussion

The chapter is where Part III closes, and several distinct strands converge here that have been building across the preceding chapters. The strategic-composition derivation lifts the symmetric partially-opposing case from Ch.4’s scalar-coupling framing to equilibrium machinery, completing the contraction-to-equilibrium hand-off that Ch.4 surfaced as a boundary. Agent opacity formalizes the dual of observation quality — #der-agent-opacity is the chapter’s strongest single segment, with rich ## Findings content — and the 16-cell adversarial-targeting decomposition closes a previously-reserved gap. The adversarial-exponent regime taxonomy ($b = 2, 3/2, \sim 1$) gives the superlinear tempo advantage its disturbance-model decomposition; observation-noise gating gives the structural condition under which the superlinear advantage collapses to roughly linear. Per-dimension persistence — formally located here but content-wise belonging with Part I Ch.4 — gives adversarial concentration its mathematical form. The chapter delivers the framework’s first operational instance of identifiability-floor pattern Instance 3 (composite-contraction certification from component-marginal data), and the closing notes hand off cleanly to TST and logogenic-agents work.

Strategic composition lifts contraction to equilibrium.

#deriv-strategic-composition is the chapter's central machinery move. When sub-agents are individually Class 1 (Separated — directed separation holds by construction) but pursue *partially-opposing* objectives, *the composite transitions on the dynamic-regime axis from R0 contraction-regime to R1 equilibrium-regime (under α') or R2 cyclic-distributional-regime (under β')*, per #disc-dynamic-regime-axis. The architectural classification from Part II Ch.1 is *preserved* under composition with goal divergence (under goal-blind routing and distinct sub-agent substrates, per #der-directed-separation Composite-level inheritance) — only the dynamic regime changes. The architectural class is set by the processing-pathway coupling, not by belief content (a sub-agent modeling another's goal-state is ordinary belief content, not a class-changing coupling). The framework's contribution is making the regime transition precise via equilibrium-convergence framing: under sub-scope α' (potential games — Monderer-Shapley 1996; monotone games — Rosen 1965 diagonally-strictly-concave), AAT's persistence machinery transfers cleanly with A2'-analog conditions derived; under sub-scope β' , the framework reaches only CCE set-convergence rather than last-iterate convergence (the honest scope limit). The segment also provides the formal home for *effects spiral* via the joint-Jacobian eigenvalue condition — what previously read as a metaphor for cascading misalignment becomes a derivable structural property of the composite at R1.

The implications-side payoff has three threads worth surfacing. First, the result *closes the framework's contraction-to-equilibrium hand-off* that Ch.4's #result-contraction-template named as a boundary. Saddle-point equilibria break attracting-fixed-point under Slotine 2003; passivity universality fails under adversarial inputs; Daskalakis et al. 2018's last-iterate non-convergence completes the trio of obstructions. The strategic-composition derivation supplies the *replacement* machinery for those obstructions: equilibrium theory applies in the partially-opposing case where contraction analysis fails. The framework's posture is that *scope-honesty is load-bearing architecture* — knowing where each tool applies cleanly is more useful than trying to extend one tool across all regimes, and the chapter delivers the equilibrium-theoretic tool that handles where contraction stops.

Second, the regime-transition result composes with #der-directed-separation's axis-decomposed inheritance table into a sharper structural commitment: *the framework cannot avoid equilibrium-regime / cyclic-distributional-regime treatment by limiting attention to architecturally-Class-1 sub-agents*. As soon as sub-agents have partially-opposing objectives — which is the empirical norm in multi-agent settings, including most human organizations and many ML systems — the composite operates in R1 or R2 on the dynamic-regime axis, requiring equilibrium-theoretic machinery rather than the contraction-machinery applicable in R0. The architectural-class axis and the dynamic-regime axis are independent: the wrapping construction (#der-class-coercion-via-wrapping) operates *only* on architectural class (Class 3 component $\rightarrow$ Class 1 composite via external scaffolding at structural-leakage cost); alignment work (#hyp-symbiogenic-composition symbiogenic absorption is the canonical mechanism) operates *only* on dynamic regime (R1 $\rightarrow$ R0 via objective re-alignment). The two axes interact but do not collapse — wrapping a strategic composite recovers architectural Class 1 at the wrapper level but the dynamics remain R1 (the equilibrium failure modes — multi-equilibria, saddle-Nash, equilibrium selection ambiguity — persist regardless of wrapping). The framework's posture is that the dynamic regime is a *property of the game structure* (objective alignment + potential/monotone structure + population scale), and goal divergence is one of two operations (alongside adversarial coupling pressure from #disc-adversarial-coupling-pressure) that move composites along their respective parallel-pattern axes — modularity-state on Axis A architectural class; dynamic regime on Axis B per #disc-dynamic-regime-axis §"Cross-axis interactions."

Third, the chapter formalizes mechanism-design impossibility (Gibbard-Satterthwaite 1973-75 / Myerson-Satterthwaite 1983 / Arrow 1951) as the *designer-side sister meta-pattern* #disc-implementation-impossibility to the agent-side #disc-identifiability-floor. The three classical theorems land as the cluster's charter instances (#deriv-strategy-proofness-impossibility, #deriv-bilateral-trade-impossibility, #deriv-social-welfare-aggregation-impossibility), with sub-scope α' of #deriv-strategic-composition named in two of them (GS and Arrow) as *adjacent to but not identical with* the preference-domain-restriction escape (single-peaked / single-crossing) — convergence-under-current-reports is not dominance-over-all-reports. The 2026-05-20 strengthen-first arm's three-reframing test (see #deriv-strategy-proofness-impossibility Discussion) is the documented argument for the non-identity. The agent-side / designer-side

suffix discipline at meta-segment naming (-floor / -impossibility) tracks the actor-positioning distinction in the canon; the cluster's discipline filter is constructive-impossibility-shape only, deliberately excluding constructive mechanism design (VCG, revenue equivalence).

Symbiogenic composition and the asymmetric absorption mechanism.

#hyp-symbiogenic-composition (formally in Ch.1) gives the chapter its account of *how composite agents form* from sub-agents that were previously not composed: asymmetric absorption, where a host integrates an endosymbiont and the resulting composite's teleological unity U_O crosses the scope threshold from below. The implication is concrete: the framework recovers the empirical pattern of biological symbiogenesis (mitochondria, chloroplasts, gut microbiota-host integration) and organizational acquisition (objective absorption, function transfer, autonomy reduction) as the *asymmetric-parameter limit* of #deriv-critical-mass-composition's closed form. Specifically, when sub-agent parameters $(\alpha_A, R_A, \gamma_A)$ and $(\alpha_B, R_B, \gamma_B)$ are *strongly asymmetric* (one sub-agent dominates), the composite sector constant α_c collapses toward the dominant sub-agent's α , and the absorbed sub-agent's objectives are progressively folded into the dominant sub-agent's via the U_O multiplier on γ . The framework's posture is that symbiogenesis is *not a separate kind of composition* — it is the asymmetric-parameter limit of the same composition machinery that handles the matched-Tier dyad and the heterogeneous-Tier cases.

The cross-domain transfer is concrete enough to ground in specific examples. The eukaryotic-cell formation (Margulis's serial endosymbiosis) is the canonical biological instance: an archaean host absorbing a bacterial endosymbiont produces a composite with the host's broader objective surface and the endosymbiont's specialized metabolic function. Corporate acquisition produces the same dynamics at a different scale: the acquirer's strategic surface absorbs the target's operational function, with the target's objectives folded in through the U_O alignment process. The framework's contribution is the *mathematical form* — the closed-form critical-mass derivation specialized to the asymmetric-parameter limit — that makes the symbiogenic process predictable rather than purely empirical.

Modular safety architectures fail under goal divergence.

The most consequential implication of strategic composition for AI-safety practice is what #28 in the catalog names but does not yet have a segment home: modular safety architectures of the form “compose modular safety modules with a central planner” are *structurally guaranteed to fail under goal divergence*. The argument composes #deriv-strategic-composition (Class-1 sub-agents with partially-opposing objectives transition the composite from R0 contraction-regime to R1 equilibrium-regime or R2 cyclic-distributional-regime per #disc-dynamic-regime-axis) directly: when safety modules and the central planner have non-aligned objectives — exactly when the safety modules are *needed* (to constrain the planner against its own objectives that may produce undesired behavior) — the composite operates in R1 or R2 on the dynamic-regime axis, and the modular safety guarantees of the individual safety modules (designed for R0 contraction-regime dynamics with their unique-attracting-fixed-point machinery) do not transfer to the composite's equilibrium-regime or cyclic-distributional-regime dynamics. The failure mechanism is dynamic-regime change, *not* architectural-class change (the composite remains architecturally Class 1 under goal-blind routing per #der-directed-separation Composite-level inheritance; what fails is the *regime-dependent* machinery that the safety guarantees assumed).

The framework's reading of several empirical AI-safety patterns follows directly. Red-team penetrations of constitutional AI exploit the structural mismatch between the constitutional constraints (one set of objectives) and the model's training objectives (another set); under partial opposition the composite operates in R1 equilibrium-regime or R2 cyclic-distributional-regime — admitting saddle-Nash equilibria, multi-equilibria with distinct basins, and last-iterate non-convergence (Daskalakis et al. 2018) — regardless of each component's nominal modularity. Mesa-optimizer formation within nominally-modular systems (Hubinger et al. 2019) instantiates the same regime-axis failure: a sub-component developing its own objective surface that opposes the outer system's objective produces partial opposition that shifts the composite from R0 contraction-regime to R1 equilibrium-regime, and the composite's modular guarantees (designed for R0 contraction-machinery) collapse at the regime boundary. The recurring difficulty of “scalable

oversight” without per-module identifiability is the chapter’s identifiability-floor Instance 3 instance manifesting: composite contraction certification from component-marginal data is structurally impossible (Liberzon 2003), and oversight regimes that rely on component-only inspection cannot, in general, certify the composite’s safety properties.

The framework’s prescription is not that modular safety architectures cannot work — it is that they require either *goal-alignment commitment by design* (sub-agents constructed to share objectives, removing the partial-opposition condition; keeps the composite in R0 contraction-regime where modular guarantees compose cleanly) or *equilibrium-theoretic treatment from the start* (acknowledging the composite operates in R1 or R2 and applying #deriv-strategic-composition’s α'/β' machinery — joint-Lyapunov-on-deviation under potential/monotone structure for R1; CCE convergence at $O(1/\sqrt{T})$ for R2 — with explicit acknowledgment of equilibrium-selection ambiguity and saddle-Nash failure modes). The implications-side payoff for AI-safety methodology is concrete: any safety architecture that *assumes* modular guarantees compose at the composite level is making a structural commitment that holds only under R0 — i.e., only under goal alignment. Where goal alignment cannot be designed-in by construction (which is the empirical case for most current deployment settings), the framework prescribes the equilibrium-theoretic machinery at the composite scope (architectural wrapping per #der-class-coercion-via-wrapping remains available as an *independent* tool operating on the orthogonal architectural-class axis — it recovers Class 1 at the wrapper level but does not change the dynamic regime; the equilibrium-failure modes persist regardless of wrapping). See #disc-dynamic-regime-axis §”Cross-axis interactions” for the two-axis treatment.

Agent opacity, the $U_o \leftrightarrow H_b$ duality, and 16-cell adversarial targeting.

#der-agent-opacity carries the chapter’s strongest single-segment finding, well-developed in its ## Findings section. *How well the agent sees the world* (U_o) and *how well the world sees the agent* (backward predictive uncertainty H_b) are formal duals. Both quantify information flow through the agent-environment boundary, in opposite directions; both are observer/horizon/trajectory-indexed; and the same H_b quantity has *opposite value to the agent* depending on coupling sign — cooperative coupling rewards low H_b (allies must predict to coordinate), adversarial coupling rewards high H_b (adversaries cannot neutralize what they cannot predict). The sign-flip is not a separate posit; it falls out of the signed-coupling structure from Ch.4.

The implications-side payoff that the segment’s Findings flag but is worth elaborating is the *16-cell joint composition* with the recipient-side classification from Ch.4. Four emitter regimes (high/low H_b , cooperative/adversarial) $\times$ four recipient regimes (informative / magnitude-shock / structural-shock / ambient-noise) produce 16 cells, each with a different optimization signature for “where to attack” or “where to coordinate.” The closed-form arg-max for adversarial targeting routes to the cell where the attacker’s opacity is high AND the recipient’s regime detector treats the event as structural-shock (which the recipient cannot absorb without architectural change). The framework therefore replaces the OODA metaphor “inside the opponent’s loop” with a precise bilinear optimization problem — the attacker chooses the cell that maximizes the recipient-side regime-induced repair cost, given the attacker’s opacity budget.

The decomposition $\mathcal{T}^{\text{eff}} = \mathcal{T} \cdot H_b / H_b^{\text{max}}$ gives the framework’s quantitative reading of the OODA-loop tempo-opacity multiplier. The superlinear $(\mathcal{T}_A / \mathcal{T}_B)^2$ adversarial scaling from Ch.4 acquires a second multiplier: the *bilateral opacity ratio* $(H_b^{A|B} / H_b^{B|A})^2$ under Model D, with the corresponding 3/2 exponent under Model S. The framework predicts that high-opacity attackers (those who can act without being predicted by the target) get *quadratic* opacity-leverage on top of *quadratic* tempo-leverage — total quartic scaling in the opacity-times-tempo regime. This is the formal grounding for what Boyd’s framing called “operating inside the OODA loop”: tempo and opacity multiply, and the multiplication is quadratic-in-quadratic at the structural extreme.

The duality also unifies what would otherwise be parallel theorems. High- U_o low- H_b agents (sees well, is seen well) sit at one corner of the duality; low- U_o high- H_b agents (sees poorly, is hidden) sit at the opposite corner. The same machinery covers both via the signed-coupling structure rather than separate analyses, and the cross-domain transfer to multi-agent safety (legibility for cooperation), adversarial robustness (opacity as targeting-vulnerability shaping), and IDT-style monitoring of opaque AI systems all follow from the same dual.

Adversarial scaling regimes and the OODA self-bound.

#result-adversarial-exponent-regimes and #obs-gated-tempo-advantage together decompose the superlinear-tempo-advantage from Ch.4 into a regime-typed taxonomy: $b = 2$ under deterministic drift (coupling-dominant); $b = 3/2$ under stochastic noise (coupling-dominant); $b \rightarrow 1$ in the coupling-non-dominant regime where observation noise is high enough to gate the tempo advantage to near-linear scaling. Both segments are AAT-internal derivations; neither carries a ## Findings section, so the catalog-grade treatment lives in this implications segment.

The implication worth surfacing is that the *non-superlinear regime* is operationally important and predictable from AAT's existing machinery. When observation noise $U_{o,B}$ is high enough that the attacker's tempo advantage doesn't penetrate the defender's observation channel cleanly, the gain-principle update $\eta_B^* = U_{M,B}/(U_{M,B} + U_{o,B})$ collapses toward zero, and the defender's effective correction rate α_B shrinks. But — counterintuitively — this *helps* the defender at the system level by gating the exponent. The framework predicts that *cheap noise injection* into the defender's observation channel can be a defensive strategy against a high-tempo attacker, by reducing the exponent from superlinear to linear, even at the cost of slowing the defender's individual updates. This is the structural reason the empirical pattern that highly-noisy environments produce more cooperative-game-like dynamics than expected — observation noise gates the adversarial advantage to near-linear scaling, where the attacker cannot exploit superlinear scaling to drive the defender past the destabilization threshold.

The composition with detection latency from #deriv-update-detection-latency produces a cross-segment finding (#40 in the catalog — *OODA Tempo + Detection Latency Self-Bounds Adversarial Targeting*) that has no other natural home. The 16-cell adversarial-targeting matrix is *self-bounded*: the most damaging cell to attack is the one the defender's *detection latency* makes hardest to recognize. The composition is $\mathcal{J}^{\text{eff}} = \mathcal{J} \cdot H_b/H_b^{\text{max}}$ (from #der-agent-opacity) and $\mathbb{E}[T_{\text{detect}}] = \Omega((n+1)/\varepsilon)$ (from #deriv-update-detection-latency): the cell with maximal damage is the one with maximal opacity and minimal ε given the defender's $n_{\min}$. Adversarial targeting therefore concentrates *where defender's response is slowest*, and the framework predicts the empirical pattern that cybersecurity incidents cluster on long-trusted dependencies (longest accumulated $n_{\min}$) and that adversarial-ML examples cluster on high-confidence model regions. Both observed empirically; the framework supplies the unifying mechanism.

The cross-segment finding's prescription is concrete: defenders should *short-circuit accumulated trust* on critical dependencies (forced forgetting to reset $n_{\min}$) and *probe high-confidence model regions* with explicit interventional checks. The detection-latency self-bound says these are not optional optimizations — they are the structural countermeasures the adversary's optimal-targeting argument forces the defender to deploy. The framework's contribution against the empirical adversarial-robustness literature is that the targeting concentration has a structural derivation, not merely a folkloric observation, and the countermeasures follow directly.

Closing Part III.

Per-dimension persistence (#result-per-dimension-persistence) is formally in this chapter though its content belongs with Part I Ch.4. The implications-segment treatment landed at Part I Ch.4; here it appears as the adversarial concentration mechanism: an opponent who identifies the target's weak dimension can concentrate disturbance there, amplifying the mismatch ratio asymmetrically while the aggregate L^2 metric still looks acceptable. The framework's prescription — *per-dimension monitoring is a structural requirement for any agent operating under directed pressure* — bites hardest here, where strategic composition delivers attackers who can locate the weak direction and adversarial targeting decomposes the attack into recipient-side regimes that exploit the dimension-asymmetry.

The chapter is also where Part III's open population-dynamics questions surface as recognized gaps: latent structural diversity (variation invisible to current performance but consequential under regime change), endogenous coupling (γ as function of population composition rather than exogenous parameter), composition-transition dynamics (Miller 2022's extreme transition motif), and computational thresholds for social behavior (Miller's ICE framework). These four gaps are the bridge to *population-level* dynamics that the current pairwise/strategic-equilibrium scope does not

cover. When fleshed out they would likely form their own chapter on population dynamics, possibly relocating the per-dimension result alongside. The implications-segment treatment flags this rather than over-claiming present coverage.

What Part III closes on is the composite-agent machinery in its operational form: closure-defect / bridge-lemma representability criterion (Ch.2), unity dimensions and Auftragstaktik bandwidth allocation (Ch.3), signed-coupling structure with cooperative-vs-adversarial as one inequality and the recipient-side four-regime taxonomy (Ch.4), strategic composition under equilibrium machinery and agent opacity with 16-cell adversarial targeting (Ch.5). The four chapters compose into AAT's full composite-agent treatment, with the identifiability-floor pattern's four formal instances (Instances 1–4) covering the framework's scope-honest commitments at the boundaries.

The bridge into TST (02-tst-core/) and the logogenic-agent stages (03-llm-core/, 04-eli-core/) is direct from here. TST instantiates the AAT machinery in the software domain, with code quality as observation infrastructure, Lindy-Effect as the temporal-optimality grounding, and software development as the *privileged calibration laboratory* (high P1–P6 identification properties). Logogenic agents — Class 3 (Coupled) by construction — inherit the wrapping construction from Ch.2 and the bias bound from Part II Ch.5, with the framework's reach into LLM-class substrates running through those two segments. The orient-cascade-recovery-via-scaffolding argument from Part II Ch.5 implications is the structural claim that connects AAT's persistence machinery to production agentic systems.

The implications-segment series for AAT chapter-ends is complete with this segment. The TST and logogenic chapter-ends follow the same pattern; that work picks up under task 11 after Joseph's check-in.

Working Notes.

- **Migration note (2026-05-09 GUC rename; residue fix 2026-05-21):** Class 2 ↔ Class 3 swap. Pre-2026-05-09: Class 2 = fully merged, Class 3 = partially modular. Post: Class 2 = Partial, Class 3 = Coupled. The canonical strategic-composition pattern “Class-1-sub-agents → Class-3-composite” (old vocab; Class 3 = partially modular) is now “Class 1 (Separated) sub-agents → Class 2 (Partial) composite.” This segment carried the pre-rename “Class 3” label through the 2026-05-09 sweep without swapping; the residue was identified and the literal-label swap landed 2026-05-21 per spikes/strategic-composition-class-3-attempt-2026-05-21/99-VERDICT.md (F4) and 07-SOLIDIFIED-PLAN.md Phase 1 (commit 0884e12). The deeper rewrite — replacing architectural-class-change mechanism with dynamic-regime-change mechanism in the §”Strategic composition lifts contraction to equilibrium” and §”Modular safety architectures fail under goal divergence” implications — landed 2026-05-21 per 07-SOLIDIFIED-PLAN.md Phase 4 (4c); the mechanism is now dynamic-regime transition (R0 → R1 / R2 per #disc-dynamic-regime-axis) with architectural class preserved under goal-blind routing (per #der-directed-separation axis-decomposed inheritance table). Removed at candidate stage per FORMAT.md Gate 4.
- This segment is a chapter-end discussion grouping findings from #deriv-strategic-composition, #der-agent-opacity, #result-adversarial-exponent-regimes, #obs-gated-tempo-advantage, #result-per-dimension-persistence (content cross-referenced from Part I Ch.4 implications), #hyp-symbiogenic-composition (formally in Ch.1, content surfacing here), and #der-interaction-channel-classification (Ch.4, the recipient-side dual to this chapter's emitter-side classification). Cross-pattern reading leans on #disc-identifiability-floor for Instance 3's operational form, #deriv-update-detection-latency for the OODA self-bound composition (#40), and #disc-modularity-state-dynamics for the goal-divergence-as-modularity-decreasing operation.
- **Cross-reference vs canonical-home policy.** Two source segments in this chapter carry well-written ## Findings sections — #der-agent-opacity and #result-per-dimension-persistence. Both are cross-referenced rather than restated. The findings whose canonical home is this implications segment are: #23 (Class-1-sub-agents under partially-opposing objectives — composite transitions from R0 contraction-regime to R1 equilibrium-regime / R2 cyclic-distributional-regime per #disc-dynamic-regime-axis; architectural class is preserved, dynamic regime is what changes), #28 (modular safety architectures fail under goal divergence — no segment home), #44 partial (the exponent-regime decomposition $b = 2, 3/2, \sim 1$), #54 (observation noise gates advantage), #52 (symbiogenic composition dynamics, asymmetric-limit reading), #40 (OODA tempo + detection latency self-bounds — cross-segment finding without home), and #50 (strategic composition game-theory bridge — companion to #23).

- **#28 Modular Safety as segment-promotion candidate.** The catalog draft entry #28 references a finding that has no home segment. The implications segment is its canonical home for now; a future cycle may promote a dedicated #disc-modular-safety-under-goal-divergence segment under Appendix A with ## Findings, at which point this segment downgrades to cross-reference.
- **Cross-segment finding without home — #40 OODA + Detection Latency Self-Bound.** Composes #der-agent-opacity (16-cell), #result-adversarial-tempo-advantage (superlinear), and #deriv-update-detection-latency ($\Omega((n + 1)/\epsilon)$) into the self-bounded adversarial-targeting prediction. No source segment carries this composition; the implications segment is its canonical home. Empirical signature (cybersecurity incident concentration on long-trusted dependencies, adversarial-ML clustering on high-confidence regions) is observed but not pulled together; the framework's contribution is the unifying mechanism.
- **Identifiability-floor pattern's full four-instance treatment + designer-side sister meta-pattern.** The agent-side identifiability-floor cluster is now fully surfaced in the implications-segment series: Instance 1 (Part II Ch.4-strategy-dynamics — on-policy L0-insufficiency detection), Instance 2 (Part II Ch.3-strategy-structure — L1' mixture identifiability under unobservable common cause), Instance 3 (Part III Ch.2-composition-machinery, operationalized here — composite contraction certification from component marginals), Instance 4 (Part II Ch.5-orient-cascade / Appendix A — architecture-noidentifiability from on-policy summary data, #der-architecture-noidentifiability, dual-anchored on Kalman-Ho canonical-form non-uniqueness and CHT-at-agent-as-SCM). Instance 5 (Mehra non-identifiability in adaptive Kalman) remains candidate. Mechanism-design impossibility (Gibbard-Satterthwaite / Myerson-Satterthwaite / Arrow) tested as a candidate fifth instance via the strategic-composition derivation in this chapter and confirmed *not* to fit the identifiability-floor's actor-positioning (designer-side vs agent-side; the 2026-05-20 strengthen-first arm's record); it lands instead as the three charter instances of the designer-side sister meta-pattern #disc-implementation-impossibility, alongside the agent-side value-functional-grounding meta-pattern #disc-value-functional-grounding-floor (Result G' + Cohen-2022 strengthened). The three-cluster constructive-impossibility-posture taxonomy (#disc-constructive-impossibility-posture) is fully realized end-to-end.
- **Connection to NeurIPS submissions.** The strategic-composition machinery is closest to a future multi-agent extension of Paper 2's convergence theory; the modular-safety implication is closest to a future AI-safety-track submission. Current cross-references in Paper 2 / Paper 3 are scope and structural-argument links rather than theorem-grade adoption from this chapter.
- **AAT chapter-end implications series complete.** With this segment, the AAT-core chapter-end implications work covers Part I Ch.4 + Part II Ch.2/3/4/5 + Part III Ch.2/3/4/5 — seven chapter-end segments plus the pilot. Each cross-references existing ## Findings where they exist and treats canonical-home findings in full where they do not. Cross-segment findings (CS1 unified RL convergence, CS2 compounding stability-induced myopia, CS3 three matrix-form bounds on Fisher backbone, #40 OODA self-bound) land in the chapter whose last prerequisite completes. The optional Appendix-A end segment (CS3 + cross-references to M1/M2/M3/M4) and the TST / logogenic-agents chapter-ends are remaining work flagged in tasks 10 and 11.

Incidental audit gold (gold-lift sweep, A15, 2026-05-31). *Orthogonal* chapter-level synthesis / framing material staged for an eventual separate promotion pass. **Coverage:** no auditor produced a dedicated reflection on this implications segment (the de-novo walks stopped at #der-agent-opacity as the last core Part III segment, or treated the chapter at the level of its constituent results); the only impl-* reflection in the contributing dirs is 526815's 92-impl-cooperative-adversarial, which belongs to the *prior* chapter (A14), not here. The chapter-level synthesis gold below is harvested from the §III-completion reflections; the per-result gold (squared-law, fog-of-war, min()-not-average(), $U_o \leftrightarrow H_b$ duality, etc.) was routed to the constituent segments #deriv-strategic-composition / #der-agent-opacity / #result-adversarial-exponent-regimes / #obs-gated-tempo-advantage / #result-per-dimension-persistence, which is the correct home per the match-by-content discipline.

2. Candidate Discussion.

- **The chapter as the control-theory $\leftrightarrow$ game-theory weld.** Multiple auditors named this chapter the place where “Section I's survival bounds and game theory's convergence proofs become the same physics” — the sector-persistence template (single-agent survival) transferring to the equilibrium layer (multi-agent convergence) via α_{joint} , with H_b supplying the information-flow boundary condition. A candidate chapter-synthesis framing for the implications opener: Part III's adversarial/strategic

results are not a separate theory bolted on, but the *same* persistence machinery read across the agent-environment boundary in both directions (Gemini 849201, Claude 471203, Gemini 193847 — convergent).

Belongs elsewhere.

- **Whole-framework synthesis (README / OUTLINE preamble, 04-eli-core/ reach).** One auditor's end-of-audit synthesis is framing-level rather than segment-level: AAT “unites control theory (Lyapunov), information theory (IB/MDL), causal inference (Pearl/do-calculus), and organizational sociology (Brooks's Law, OKRs, symbiogenesis) under one coherent mathematical umbrella,” with the primary *weakness* being integration debt (editorial bloat in src/ Findings blocks, topological forward-references, residual TF-06/TF-11 archaeology) rather than the math — “a diamond that needs its scaffolding removed and edges polished” (Gemini, 829314). This is README/positioning material and an integration-debt signal, not implications-segment content; recorded here because it surfaced at the chapter that closes the AAT-core walk. (*The integration-debt items are certified-track housekeeping, not gold — see off-ramp.*)

Off-ramp (NOT gold) — routed for adjudication, not promotion.

- (829314, 849201 — recurring across Part III) — “editorial bloat”: in-src ## Findings blocks read as misplaced. Two Gemini auditors repeatedly flagged the large ## Findings / search-log / related-work blocks pasted into Part III result segments as “destroying the pedagogical flow,” recommending relocation to an appendix or tracking system. **This is a false alarm about project mechanics** (the in-segment ## Findings block is by-design the source the auto-extracted root FINDINGS.md is generated from), so it is *not* an action — but it is preserved as the same de-novo-reader stumble two fresh readers independently hit (also recorded on #result-persistence-condition in the pilot lift), a mild signal that the convention is not self-evident. No segment change.
- **Integration-debt housekeeping** (829314 final synthesis): topological forward-references in Part III (e.g. #der-agent-opacity referencing #der-interaction-channel-classification before it appears in a linear read), and residual TF-06/TF-11 archaeology in footers — routed to the integration-cleanup / OUTLINE-ordering stream, not theory. Out of scope for this WN-only lift; the C-iv idiom-drift residue in the OUTLINE rows (lines 91, 323, 328) is the most concrete instance and is flagged on #deriv-strategic-composition's off-ramp.

[Gap] Latent structural diversity: variation in correction architectures invisible to persistence analysis, consequential under regime change

[Gap] Endogenous coupling: γ as function of population composition, not exogenous parameter; coupling emergence threshold

[Gap] Composition transition dynamics: epochal stability $\rightarrow$ latent diversification $\rightarrow$ niche emergence $\rightarrow$ cascading restructuring $\rightarrow$ re-equilibration (adopts Miller 2022's extreme transition motif)

[Gap] Computational thresholds for social behavior: minimum agent complexity and interaction depth for composition dynamics (adopts Miller 2022's ICE framework; grounds 9.9)

Part IV

Appendices: Details

Appendix A

Derivation: Sector Condition Stability — Lyapunov Derivation

The Lyapunov machinery underlying the sector-condition stability result and the central persistence inequality. The mismatch dynamics are taken in general nonlinear vector form, $\dot{\delta} = -F(\mathcal{T}, \delta) + w(t)$, with the correction function F required to satisfy only the *qualitative* properties (A1)–(A3) and the local sector condition (A2') stated in `#form-sector-condition`. The Lyapunov function is the canonical quadratic $V = \frac{1}{2}\|\delta\|^2$; the steady-state results follow from standard Lyapunov / Itô-Lyapunov machinery applied to AAT's correction-function object.

The derivation produces the framework's central persistence results in their precise form. **Proposition A.1 (Model D — bounded disturbance)**: under $\|w(t)\| \leq \rho$ and the sector condition, mismatch is ultimately bounded by $R^* = \rho/\alpha$ — exactly the persistence threshold the chapter cites, and $\mathcal{B}_R$ is positively invariant when $\alpha R > \rho$. **Lemma A.1N (threshold tightness)**: the condition is unimprovable given A2' — necessary and sufficient for the *uniform class-level* guarantee (witness $F = \alpha\delta$ under outward disturbance), and agent-level necessary exactly when the sector bound is radially tight (the linear case); for a general fixed F the only-if direction is *false* (demonstrated by the dip counterexample in the Remark following A.1N — the sector floor can sit at radii the dynamics render unreachable). **Proposition A.1S (Model S — additive stochastic disturbance)**: the region-aware Itô-Lyapunov analysis gives the stopped second-moment bound, the steady-state RMS $R_S^* = \sigma_w \sqrt{n/(2\alpha)}$ (the $1/\sqrt{\alpha}$ scaling that distinguishes Model S from Model D's $1/\alpha$ scaling), and the stationary-sharp fixed-time tail $P(\|\delta(t)\| > R) \leq n\sigma_w^2/(2\alpha R^2)$. **Proposition A.2 (adaptive reserve)**: $\Delta\rho^* = \alpha R - \rho$ — the additional disturbance an agent can absorb before mismatch leaves the sector-condition region.

The segment's load-bearing new exact result is **Corollary A.1S.1 — the disturbance-model containment dichotomy**. The persistence-region first-exit probability is *categorical*: $P(\tau_R < \infty) = 0$ under Model D (deterministic positive invariance), $P(\tau_R < \infty) = 1$ under Model S (a.s. exit of a non-degenerate diffusion from a bounded region, for *any* F satisfying A2'). The achievable value is exactly the two-point set $\{0, 1\}$, and which point obtains is fixed by the disturbance model's support structure (bounded vs. unbounded), **not** by correction strength α . Increasing α tightens the typical scale and the fixed-time tail; it cannot interpolate between the two regimes. Pathwise containment is categorically Model-D-only — additive stochastic forcing removes the *kind* of guarantee available, not merely its rate. The companion no-go demonstration (that the natural Ville/Doob maximal-inequality route to a $P(\tau_R < \infty) < 1$ bound *cannot* exist) is in `#deriv-stochastic-non-exit`. The dichotomy sharpens the hand-off into `#result-structural-adaptation-necessity`: in any genuinely stochastic environment, leaving the parametric-correction region is a certain eventual event, so structural adaptation is *generic, not exceptional*, for any sufficiently long-lived agent. The sub-scope α / β partition (where A2' is structurally derived versus where it is per-system assumed) is held at `#form-sector-condition`; the Lyapunov proofs below apply uniformly across both.

Motivation.

3.9 hypothesizes the linear ODE $d\|\delta\|/dt = -\mathcal{T}\|\delta\| + \rho$ as a first-order approximation. The linear form yields clean closed-form results but commits to a specific functional relationship between mismatch magnitude and correction rate. True correction dynamics are almost certainly nonlinear — exhibiting saturation at large mismatch, threshold effects near zero, and structural breakdown when the model class is exhausted.

A Lyapunov approach proves persistence and stability under much weaker assumptions: any correction dynamics satisfying qualitative monotonicity properties (the sector condition). The results below are strictly more general — the linear case is recovered where the sector bounds coincide.

Setup.

Let $\delta(t) \in \mathbb{R}^n$ be the mismatch vector — the difference between the model's predictions and reality across n observable dimensions. The vector treatment connects to per-dimension tempo analysis (16.5).

The mismatch dynamics are:

Definition (Dynamics Setup).

$$\frac{d\delta}{dt} = -F(\mathcal{T}, \delta) + w(t) \quad (\text{A.1})$$

where:

- $F(\mathcal{T}, \delta) : \mathbb{R}_+ \times \mathbb{R}^n \rightarrow \mathbb{R}^n$ is the **correction function** — how the agent's adaptive process reduces mismatch. It maps to the same space as δ (so that the inner product $\delta^T F$ in the sector condition is well-defined). This subsumes the update gain η^* (3.6), event rate ν , and the structure of the update rule.
- $w(t)$ is the **disturbance** — new mismatch introduced by environmental change, with $\|w(t)\| \leq \rho$ (bounded disturbance rate).

The linear case from 3.9 has $F(\mathcal{T}, \delta) = \mathcal{T} \cdot \delta$.

Assumptions on F .

The Lyapunov derivations below operate on the structural-property triple (A1) zero-correction-at-zero-mismatch, (A2') local sector condition, (A3) tempo-monotonicity, stated in full at 4.2. The sub-scope α / β partition — the agent classes where A2' is structurally derived (Bayesian / exponential family / strongly-convex gradient / L2-regularized / linear-PD) versus those where it stands as a well-scoped per-system empirical claim (PID, rule-based, human judgment, severely misspecified, variational, non-convex beyond basin, per-step SGD) — is also held there, together with the operator-family classification (Rockafellar / Bauschke-Combettes / Baillon-Haddad) that names how the AAT-internal sub-scope partition lines up with monotone-operator theory.

The Lyapunov proofs below apply uniformly across both sub-scopes — they operate downstream of A2' regardless of whether it is structurally derived or per-system assumed.

Candidate Lyapunov Function.

Definition (Lyapunov Candidate).

$$V(\delta) = \frac{1}{2} \|\delta\|^2 \quad (\text{A.2})$$

Positive definite, radially unbounded, continuously differentiable. Its level sets $V = c$ are spheres of radius $\sqrt{2c}$ in mismatch space.

Proposition A.1: Bounded Mismatch (Single Agent).

Statement. Under (A1), (A2'), (A3), with bounded disturbance $\|w(t)\| \leq \rho$, the mismatch $\delta(t)$ is **ultimately bounded**: there exists $R^* > 0$ such that $\|\delta(t)\| \leq R^*$ for all sufficiently large t , provided $R^* < R$ (the ultimately bounded region fits within the sector-condition region).

Proof.

Compute $\dot{V}$ along trajectories:

Derived (Proof Step).

$$\dot{V} = \delta^T \dot{\delta} = \delta^T [-F(\mathcal{J}, \delta) + w(t)] \quad (\text{A.3})$$

Derived (Proof Step).

$$= -\delta^T F(\mathcal{J}, \delta) + \delta^T w(t) \quad (\text{A.4})$$

By (A2'): $\delta^T F(\mathcal{J}, \delta) \geq \alpha \|\delta\|^2$

By Cauchy-Schwarz: $\delta^T w(t) \leq \|\delta\| \cdot \|w(t)\| \leq \rho \|\delta\|$

Therefore:

Derived (Proof Step).

$$\dot{V} \leq -\alpha \|\delta\|^2 + \rho \|\delta\| = -\|\delta\|(\alpha \|\delta\| - \rho) \quad (\text{A.5})$$

$\dot{V} < 0$ whenever $\|\delta\| > \rho/\alpha$ and $\|\delta\| \leq R$ (where A2' holds).

Define $R^* = \rho/\alpha$.

Invariance of $\mathcal{B}_R$. When $R^* < R$ (the persistence condition), the ball $\mathcal{B}_R$ is *positively invariant*: any trajectory starting in $\mathcal{B}_R$ remains in $\mathcal{B}_R$ for all future time. At the boundary $\|\delta\| = R$, the sector condition holds and $\dot{V} = -\|\delta\|(\alpha \|\delta\| - \rho) = -R(\alpha R - \rho) < 0$ (since $\alpha R > \rho$ by the persistence condition). Trajectories at the boundary point inward. Therefore $\mathcal{B}_R$ is invariant, and the sector condition applies to the entire future trajectory of any trajectory starting inside $\mathcal{B}_R$.

Ultimate boundedness. Within $\mathcal{B}_R$, the Lyapunov function is strictly decreasing for $\|\delta\| > R^*$ and may increase for $\|\delta\| < R^*$. All trajectories starting in $\mathcal{B}_R$ are ultimately bounded by R^* — they are driven into $\mathcal{B}_{R^*}$ and remain in a neighborhood of it (with possible oscillation at the boundary due to the disturbance).

Initial condition requirement. The result requires the trajectory to start inside $\mathcal{B}_R$ (or to be brought inside by some external mechanism). A trajectory starting outside $\mathcal{B}_R$ is not covered — the sector condition does not hold there, and the correction function may fail to reduce mismatch. This is precisely the structural adaptation regime of 4.8: an agent whose mismatch exceeds R has exhausted its model class capacity and needs structural change to re-enter the region where parametric correction works.

The persistence condition $\alpha > \rho/R$ (equivalently $R^* < R$) is therefore *sufficient* for persistence from within $\mathcal{B}_R$. $\square$

Lemma A.1N (threshold tightness — necessity is class-level, not agent-level). Over the hypothesis class $\mathcal{F}_\alpha = \{F : (A1), (A2'), (A3) \text{ hold on } \mathcal{B}_R \text{ with sector floor } \alpha\}$, the condition $\alpha > \rho/R$ is necessary *and* sufficient for the uniform containment guarantee (every $F \in \mathcal{F}_\alpha$ keeps every trajectory from $\mathcal{B}_R$ inside $\mathcal{B}_R$ under every $\|w(t)\| \leq \rho$).

Proof. Sufficiency is Prop A.1, uniform over the class. Necessity: the extremal member $F = \alpha\delta$ lies in $\mathcal{F}_\alpha$; under the outward radial disturbance $w = \rho\delta/\|\delta\|$ its radial dynamics obey $\frac{d}{dt}\|\delta\| = \rho - \alpha\|\delta\| \geq \rho - \alpha R > 0$ whenever $\alpha < \rho/R$, so from any nonzero start the trajectory exits $\mathcal{B}_R$ in finite time. (At $\alpha = \rho/R$ exactly, the trajectory asymptotes to R without crossing — containment holds, with ultimate bound R rather than $< R$.) $\square$

The lemma answers the question the old “iff” was reaching for — *when is the persistence condition also necessary?* — and the answer is **tightness, not linearity**. For a fixed agent whose sector bound is radially tight (additionally $\delta^T F \leq \tilde{\alpha}\|\delta\|^2$ with $\tilde{\alpha} = \alpha$; the linear case $F = \mathcal{T}\delta$ is the canonical instance), the same outward-disturbance construction forces finite-time exit, and the threshold is a genuine agent-level if-and-only-if. In general, the object necessity tracks is the *upper* sector constant $\tilde{\alpha}$, not the floor α : the Lyapunov argument bounds $\dot{V}$ from above via the floor, and forcing escape needs a bound from below, which only an upper sector bound supplies.

Remark (the floor is not the barrier — why agent-level necessity fails in general). For a fixed F , failure of $\alpha > \rho/R$ does *not* imply escape, because α is a worst-case infimum over $\mathcal{B}_R$ that the dynamics may render unreachable. Demonstration (1D; extends to $\mathbb{R}^n$ radially via $F(\delta) = g(\|\delta\|)\delta$): take $R = 1$, $\rho = 1$, and $g \geq 10$ everywhere on $(0, 1]$ except a narrow dip to $g = 0.1$ near $r = 0.5$. The tight sector floor on $\mathcal{B}_R$ is $\alpha = g(0.5) = 0.1 < \rho/R$ — the persistence condition fails tenfold — yet $F(r) > \rho$ on an annulus below the dip, so every admissible trajectory started at small mismatch is stopped there: an impenetrable barrier *below* the weak spot, which no admissible disturbance can reach. The supremum of $\|\delta(t)\|$ over all $\|w\| \leq \rho$ from $\delta(0) = 0$ is ≈ 0.1 (numerically confirmed, including against randomized and outward-feedback adversaries). The sharp agent-level escape threshold in 1D is a *mountain-pass* quantity — escape requires $\rho > \max_r F(r)$ over the radii the trajectory must traverse — not the global floor αR . Condition failure therefore means the containment *certificate* is lost (and escape is forced for the extremal class member and every radially tight agent), never that escape is certified for the agent at hand.

Interpretation. The ultimately bounded region has radius $R^* = \rho/\alpha$. In the linear case, $\alpha = \mathcal{T}$, recovering 4.5’s steady-state result $R^* = \rho/\mathcal{T}$ exactly. But Proposition A.1 holds for *any* correction function satisfying the sector condition, not just the linear one.

The persistence threshold, generalized. The persistence condition $\rho/\alpha < R$ guarantees persistence for any correction function satisfying the sector condition, and Lemma A.1N shows it is *unimprovable given only A2'*: necessary and sufficient for the guarantee at class level, and agent-level necessary exactly in the radially tight case. If the correction function breaks down (A2' fails) before R^* is reached, the agent may diverge. Condition failure IS 4.8’s trigger in certificate form: when $\rho/\alpha > R$ (the environment demands more correction than the model class can *certifiably* provide), parametric adaptation loses its containment guarantee and structural change is required to restore one.

Proposition A.2: Stability Margin (Adaptive Reserve).

Statement. Under the conditions of A.1, the agent can tolerate a sudden increase in disturbance rate of:

Derived (adaptive-reserve).

$$\Delta\rho^* = \alpha R - \rho \quad (\text{A.6})$$

without mismatch diverging (where R is the radius of the sector-condition region from A2'). Beyond this, R^* exceeds R and the correction function may fail.

Proof. After a shock, the new disturbance rate is $\rho + \Delta\rho$. The new ultimately bounded radius is $(\rho + \Delta\rho)/\alpha$. This remains within the valid region iff $(\rho + \Delta\rho)/\alpha \leq R$, i.e., $\Delta\rho \leq \alpha R - \rho$. $\square$

Interpretation. $\Delta\rho^*$ is the agent's **adaptive reserve** — how much additional environmental volatility it can absorb before its model breaks down. This is a single number characterizing an agent's robustness to shock:

- An agent operating well below capacity ($\rho \ll \alpha R$) has a large reserve — it is **robust**.
- An agent near its limit ($\rho \approx \alpha R$) has a small reserve — it is **fragile**.

Table A.1

<i>Domain</i>	<i>Large $\Delta\rho^*$ (robust)</i>	<i>Small $\Delta\rho^*$ (fragile)</i>
Control	Kalman filter on slow target	Same filter on erratic target
Biology	Organism in stable niche	Same organism under climate change
Organization	Well-capitalized firm, stable market	Startup in volatile market
Military	Force with operational depth	Force at culmination point

Proposition A.1S: Bounded Mismatch Under Stochastic Disturbance.

Statement (region-aware form). Under (A1), (A2') on $\mathcal{B}_R$, (A3), with stochastic disturbance (GA-2S: $w(t)$ is zero-mean with $\mathbb{E}[\|w(t)\|^2] = \sigma_w^2$), let $\tau_R = \inf\{t : \|\delta(t)\| > R\}$ be the first-exit time from the sector-condition region. Then:

(i) *Stopped bound* — the stopped process satisfies for all $t \geq 0$:

Derived (stochastic-bounded-mismatch, stopping-time localization, Khasminskii 2012 ch. 5).

$$\mathbb{E}[\|\delta(t \wedge \tau_R)\|^2] \leq \|\delta(0)\|^2 e^{-2\alpha t} + \frac{n\sigma_w^2}{2\alpha} \quad (\text{A.7})$$

(ii) *Mean-square persistence condition* — when $R_S^* := \sigma_w \sqrt{n/(2\alpha)} < R$, equivalently

$$\alpha > \frac{n\sigma_w^2}{2R^2} \quad (\text{A.8})$$

the ultimately-bounded RMS radius fits inside the sector-condition region.

(iii') *Fixed-time tail* — under the mean-square persistence condition (ii), Markov's inequality on the stopped second moment gives, for each fixed t (sharpest in the stationary $t \rightarrow \infty$ limit),

$$P(\|\delta(t)\| > R) \leq \frac{\mathbb{E}[\|\delta(t \wedge \tau_R)\|^2]}{R^2} \xrightarrow{t \rightarrow \infty} \frac{n\sigma_w^2}{2\alpha R^2}. \quad (\text{A.9})$$

This controls the mismatch at any single time. It is *not* an infinite-horizon containment statement: under additive Brownian forcing the diffusion is recurrent on $\mathbb{R}^n$ (the additive-noise generator has no bounded non-constant harmonic function), so $P(\tau_R < \infty) = 1$ over an unbounded horizon — there is no $P(\tau_R < \infty) < 1$ bound, and none is claimed. Pathwise containment of $\mathcal{B}_R$ is a Model-D guarantee only (Prop A.1's positive invariance); Model S controls the *typical scale* (ii) and the *fixed-time tail* (iii'), not the sample path over an unbounded horizon. The kind-of-guarantee dichotomy this exposes is itself a result — see the Discussion.

(iv) *Finite-horizon sample-path bound* — when a sup-over-an-interval statement is wanted, the sound one is the additive first-exit bound

$$P\left(\sup_{0 \leq s \leq T} \|\delta(s)\| > R\right) \leq \frac{\|\delta(0)\|^2 + n\sigma_w^2 T}{R^2}, \quad (\text{A.10})$$

rigorous under (A2') alone (Khasminskii 2012 ch. 5). It controls the whole path on $[0, T]$ — stronger than (iii') within its regime — but grows linearly in T and is vacuous for $T > rsimR^2/(n\sigma_w^2)$, consistent with $P(\tau_R < \infty) = 1$.

The proof below establishes the Grönwall-type bound first, then the stopping-time localization at the end.

Proof.

The SDE form of the mismatch dynamics under stochastic disturbance is:

$$d\delta = -F(\mathcal{J}, \delta) dt + \sigma_w dW_t \quad (\text{A.11})$$

where W_t is a standard n -dimensional Wiener process.

Apply Itô's formula to $V(\delta) = \frac{1}{2}\|\delta\|^2$:

Derived (Proof Step).

$$dV = \delta^T(-F) dt + \delta^T \sigma_w dW_t + \frac{1}{2} \sigma_w^2 n dt \quad (\text{A.12})$$

The last term is the Itô correction: $\frac{1}{2} \text{tr}(\sigma_w^2 I_n) = \frac{n}{2} \sigma_w^2$.

Taking expectations (the Itô integral $\delta^T \sigma_w dW_t$ has zero expectation):

Derived (Proof Step).

$$\frac{d}{dt} \mathbb{E}[V] = \mathbb{E}[\delta^T(-F)] + \frac{n}{2} \sigma_w^2 \quad (\text{A.13})$$

By (A2'): $\delta^T F \geq \alpha \|\delta\|^2 = 2\alpha V$. Therefore:

Derived (Proof Step).

$$\frac{d}{dt} \mathbb{E}[V] \leq -2\alpha \mathbb{E}[V] + \frac{n}{2} \sigma_w^2 \quad (\text{A.14})$$

This is a linear ODE in $\mathbb{E}[V]$ with solution:

$$\mathbb{E}[V(t)] \leq V(0) e^{-2\alpha t} + \frac{n\sigma_w^2}{4\alpha} (1 - e^{-2\alpha t}) \quad (\text{A.15})$$

Since $V = \frac{1}{2} \|\delta\|^2$, the steady-state mean-square mismatch is:

$$\mathbb{E}[\|\delta\|^2]_{ss} = \frac{n\sigma_w^2}{2\alpha} \quad (\text{A.16})$$

The RMS steady-state mismatch is:

Derived (stochastic-steady-state).

$$R_S^* = \|\delta\|_{\text{rms}} = \sigma_w \sqrt{\frac{n}{2\alpha}} \quad (\text{A.17})$$

Persistence requires $R_S^* < R$, giving $\alpha > n\sigma_w^2/(2R^2)$.

Stopping-time localization. The Itô-Lyapunov step above used A2' ($\delta^T F \geq \alpha \|\delta\|^2$) on the entire trajectory, but A2' is posited only on $\mathcal{B}_R$. Replacing t with the stopped time $t \wedge \tau_R$ makes the argument valid: on $[0, t \wedge \tau_R]$, $\delta(s) \in \mathcal{B}_R$ almost surely, so A2' applies. The Itô integral $\int_0^{t \wedge \tau_R} \delta^T \sigma_w dW_s$ remains a martingale with zero expectation (optional stopping for L^2 -bounded martingales). The stopped Grönwall bound (i) follows. For (iii'), Markov's inequality applied to the stopped second moment (i) gives the fixed-time tail directly. The *infinite-horizon* ever-exit probability is 1 under additive Brownian forcing, and no nonnegative supermartingale certifies any $P(\tau_R < \infty) < 1$ bound — the natural Ville/Doob maximal-inequality route provably cannot exist (demonstrated in #deriv-stochastic-non-exit). (iv)'s finite-horizon sup-bound follows from the same Itô-Lyapunov step without the decay term. $\square$

Interpretation. Model D (Prop A.1) gives $R^* = \rho/\alpha$, scaling as $1/\alpha$. Model S gives $R_S^* = \sigma_w \sqrt{n/(2\alpha)}$, scaling as $1/\sqrt{\alpha}$. Doubling the correction efficiency halves the deterministic steady-state mismatch but only reduces the stochastic steady-state by a factor of $\sqrt{2} \approx 1.41$. Correction is less effective against noise than against drift. This difference in scaling propagates into the adversarial exponent regimes (16.3): $b = 2$ under Model D, $b = 3/2$ under Model S.

Tail bound. At steady state, Markov's inequality gives $P(\|\delta\| > R) \leq n\sigma_w^2/(2\alpha R^2)$. The agent stays within R with probability $\geq 1 - \epsilon$ provided $\alpha \geq n\sigma_w^2/(2\epsilon R^2)$. For the linear case (Ornstein-Uhlenbeck), the stationary distribution is Gaussian and exact tail probabilities are available. The Markov bound is the general result holding under (A2') alone.

Corollary A.1S.1: Disturbance-Model Containment Dichotomy.

For a trajectory started in $\mathcal{B}_R$ under (A1), (A2') on $\mathcal{B}_R$, (A3), the persistence-region first-exit probability is **categorical**:
Derived.

$$P(\tau_R < \infty) = \begin{cases} 0 & \text{Model D (bounded } w, \|w\| \leq \rho) \text{ under } \alpha R > \rho, \\ 1 & \text{Model S (additive stochastic, } d\delta = -F dt + \sigma_w dW_t). \end{cases} \quad (\text{A.18})$$

The Model-D value is the deterministic positive invariance of Prop A.1 (boundary-inward: at $\|\delta\| = R$, $\dot{V} < 0$ when $\alpha R > \rho$). The Model-S value holds for **every** $\alpha > 0$, **every** $\sigma_w > 0$, and **every** correction function F satisfying (A2'): a non-degenerate diffusion (additive forcing $\sigma_w dW_t$, $\sigma_w > 0$) exits any bounded region in finite time almost surely, for any locally bounded drift — no finite inward correction can defeat the Brownian crossing near $\partial\mathcal{B}_R$. The Ornstein–Uhlenbeck benchmark makes the mechanism explicit (scale function $s'(u) \propto e^{\alpha u^2 / \sigma_w^2}$ unbounded, hence recurrent; Khasminskii 2012¹ ch. 3–4), but the conclusion needs only non-degeneracy of the additive noise on a bounded region, not the linear structure. That the natural maximal-inequality route to a $P(\tau_R < \infty) < 1$ bound *cannot* exist — the question “are you sure you can’t just Doob/Ville this?” — is demonstrated in #deriv-stochastic-non-exit.

The achievable first-exit probability is therefore exactly the two-point set $\{0, 1\}$, and **which point obtains is fixed by the disturbance model’s support structure (bounded vs. unbounded), not by the correction strength α** . Increasing α tightens the typical scale ($R_S^* \propto 1/\sqrt{\alpha}$) and the fixed-time tail (iii’); it cannot move $P(\tau_R < \infty)$ off 1 under Model S, and is unnecessary for the 0 under Model D. Pathwise containment of $\mathcal{B}_R$ is categorically a bounded-disturbance property: additive stochastic forcing removes the *kind* of guarantee available, not merely its rate. This sharpens the hand-off into 4.8 — in any genuinely stochastic environment region-exit is a certain eventual event, so the structural-adaptation trigger is *generic, not exceptional*, for a sufficiently long-lived agent.

Summary of Results.

Table A.2

Result	What it proves	Assumptions	Linear case recovery
A.1 (Bounded Mismatch)	$R^* = \rho/\alpha$	(A1), (A2'), bounded ρ (GA-2)	$\alpha = \mathcal{T}$ gives $R^* = \rho/\mathcal{T}$
A.1S (Stochastic Bounded Mismatch, region-aware)	$R_S^* = \sigma_w \sqrt{n/(2\alpha)}$ (stopped bound); fixed-time tail $P(\ \delta(t)\ > R) \leq n\sigma_w^2/(2\alpha R^2)$ (stationary-sharp); no infinite-horizon non-exit bound — pathwise containment is Model-D-only	(A1), (A2') on $\mathcal{B}_R$, stochastic w (GA-2S)	$\alpha = \mathcal{T}$ gives $R_S^* = \sigma_w \sqrt{n/(2\mathcal{T})}$
Cor A.1S.1 (Disturbance-Model Containment Dichotomy)	$P(\tau_R < \infty)$ is exactly $\{0, 1\}$ — 0 under Model D (positive invariance), 1 under Model S (a.s. exit of a non-degenerate diffusion); α -invariant, correction strength cannot interpolate	(A1), (A2') on $\mathcal{B}_R$; Model D needs $\alpha R > \rho$	Model D $\rightarrow$ Prop A.1 invariance at $\alpha = \mathcal{T}$; Model S $\rightarrow$ OU recurrence is the explicit instance — exact generally (any F under A2'), not linear-scoped
A.2 (Stability Margin)	$\Delta\rho^* = \alpha R - \rho$	Same as A.1	$R \rightarrow \infty$ for linear (always stable if $\mathcal{T} > 0$)

¹Khasminskii, R. (2012). *Stochastic Stability of Differential Equations* (2nd ed.). Springer. Chapter 5: Lyapunov functions and stochastic stability; stopping-time localization for diffusion processes.

What Is Derived vs. What Is Chosen.

Table A.3

<i>Property</i>	<i>Source</i>	<i>Strength</i>
Dynamics setup $\dot{\delta} = -F(\mathcal{J}, \delta) + w(t)$ with F as correction function and w as disturbance	Definitional scope of the appendix; generalizes the linear hypothesis of 3.9	Definition
(A1) $F(\mathcal{J}, 0) = 0$	Qualitative property of any correction process by construction	Assumption (uncontroversial)
(A2') Local sector condition $\delta^T F \geq \alpha \ \delta\ ^2$ on $\mathcal{B}_R$	Sector-condition framework (Lur'e 1957); sub-scope α (Kalman/conjugate Bayesian, exponential family in natural params, gradient descent on strongly convex losses, L2-regularized convex, linear with PD KH) derived via 4.3 Prop B.3 under B1 directional fidelity (then $\alpha = \eta^* \cdot c_{\min}$); sub-scope β (PID, rule-based, human judgment, severely misspecified, variational approximations, non-convex beyond basin, per-step SGD) requires independent verification	Derived in sub-scope α; Assumption in sub-scope β (sub-scopes named explicitly in the Grounding paragraphs)
(A3) Tempo-monotonicity of $\delta^T F$	Qualitative requirement tying $\mathcal{J}$ to correction power	Assumption
Quadratic Lyapunov candidate $V = \frac{1}{2} \ \delta\ ^2$	Canonical choice for norm-bounded stability in Euclidean state spaces (Khalil 2002 ch. 4)	Formulation choice
Prop A.1: ultimate bound $R^* = \rho/\alpha$	$\dot{V} \leq -\ \delta\ (\alpha\ \delta\ - \rho)$ via A2' + Cauchy-Schwarz; standard Lyapunov ultimate-boundedness theorem (Khalil 2002 Thm 4.18)	Proved (conditional on $R^* < R$)
Positive invariance of $\mathcal{B}_R$ under the persistence condition	Boundary-inward argument: at $\ \delta\ = R$, $\dot{V} < 0$ when $\alpha R > \rho$	Proved
Lemma A.1N: threshold tightness — $\alpha > \rho/R$ is necessary-and-sufficient for the <i>uniform class-level</i> containment guarantee; agent-level necessity holds exactly under a radially tight sector ($\tilde{\alpha} = \alpha$, e.g. linear); false as an agent-level only-if for general F	Extremal-member escape ($F = \alpha\delta$, $w = \rho\delta/\ \delta\ $) + the dip counterexample (floor attained at unreachable radii; mountain-pass barrier analysis)	Proved — necessity is class-level; no fixed-agent only-if is claimed anywhere
Initial-condition scope (trajectory must start inside $\mathcal{B}_R$)	Direct consequence of A2' being local; trajectories outside are not covered	Derived (scope statement)
Identification of structural-adaptation trigger ($\rho/\alpha > R$ forces exit from parametric regime)	Connects A.1's persistence threshold to 4.8's regime boundary	Derived
Prop A.2: adaptive reserve $\Delta\rho^* = \alpha R - \rho$	Algebraic corollary of Prop A.1 under shock $\rho \rightarrow \rho + \Delta\rho$; requires R fixed under shock	Proved (corollary of A.1)
Prop A.1S (region-aware): stopped bound $\mathbb{E}[\ \delta(t \wedge \tau_R)\ ^2] \leq \ \delta(0)\ ^2 e^{-2\alpha t} + n\sigma_w^2/(2\alpha)$ + fixed-time tail $P(\ \delta(t)\ > R) \leq n\sigma_w^2/(2\alpha R^2)$ (stationary-sharp) + finite-horizon sup-bound (iv)	Itô's formula on V + A2' on $\mathcal{B}_R$ + Grönwall + stopping-time localization at τ_R (Khasminskii 2012 ch. 5; Khalil 2002 ch. 9)	Proved ; the infinite-horizon non-exit probability is structurally 1 for additive Model S (no-go, Khasminskii ch. 3–4) — pathwise containment is Model-D-only
Cor A.1S.1: containment dichotomy $P(\tau_R < \infty) \in \{0, 1\}$, α -invariant	Prop A.1 positive invariance (Model D = 0) + a.s. finite exit of a non-degenerate diffusion from a bounded region for any F under A2' (Model S = 1); Khasminskii 2012 ch. 3–4	Proved — new exact result
Stochastic steady-state RMS $R_S^* = \sigma_w \sqrt{n/(2\alpha)}$	Steady state of the Grönwall bound; scales as $1/\sqrt{\alpha}$ versus $1/\alpha$ deterministic	Derived
Tail bound $P(\ \delta\ > R) \leq n\sigma_w^2/(2\alpha R^2)$	Markov's inequality applied to steady-state second moment	Proved

Property	Source	Strength
Linear-case recovery ($\alpha = \mathcal{T}$ reproduces 4.5)	Substitution $F = \mathcal{T}\delta$ into A2'	Derived
Adversarial-exponent scaling $b = 2$ (Model D) vs $b = 3/2$ (Model S) into 16.3	Direct transfer of the $1/\alpha$ vs $1/\sqrt{\alpha}$ scalings	Derived (transfer claim)
Bounded-disturbance model (GA-2) vs stochastic-disturbance model (GA-2S) as distinct environmental regimes, not approximations to each other	Discussion-section positioning; neither regime handles heavy tails	Discussion-grade
Global sector condition would give $\Delta\rho^* = \infty$	Limit analysis of A.2 as $R \rightarrow \infty$; rejected as unrealistic for finite model classes	Derived (scope-limit observation)

The dividing line: the two persistence results (ultimate bound, adaptive reserve) and the stochastic mean-square bound are **proved** within standard Lyapunov / Itô-Lyapunov theory — the appendix's contribution is the application to AAT's correction-function object, not the mathematics. What is **chosen** is the quadratic Lyapunov candidate $V = \frac{1}{2}\|\delta\|^2$ (canonical for norm-bounded stability but not uniquely forced; see Path 3 note below) and the bounded-vs-stochastic disturbance partition (two structurally distinct empirical regimes rather than a hierarchy). The **sector condition itself** (A2') now carries its sub-scope explicitly: **derived in sub-scope** α (optimal Bayesian / exponential-family / strongly-convex-gradient / L2-regularized / linear-PD corrections, via 4.3 Prop B.3 under B1), **assumed in sub-scope** β (PID / rule-based / human-judgment / severely-misspecified / variational / non-convex-beyond-basin / per-step SGD). Prop A.1S's region condition is no longer deferred to Epistemic Status — the proposition statement is region-aware via a standard stopping-time localization (Khasminskii 2012 ch. 5); the fixed-time tail (iii') quantifies the typical-scale cost of Wiener excursions, and the *absence* of any infinite-horizon non-exit bound (a structural no-go for additive Model S) is itself the result — pathwise containment is a Model-D guarantee only.

Epistemic Status.

The setup and assumptions are *definitions* — they specify what we mean by “correction function” and “disturbance.” Propositions A.1 and A.2 are *exact* — they follow from the assumptions via standard Lyapunov theory (Khalil 2002², Chapters 4 and 9). The persistence threshold's necessity is *class-level* (Lemma A.1N, exact): $\alpha > \rho/R$ is necessary and sufficient for the uniform guarantee over the sector class, and agent-level necessary only under a radially tight sector; for a fixed general F the only-if direction is false (the demonstrated dip counterexample), so every statement of the threshold in this corpus is sufficiency-plus-class-tightness, never a fixed-agent iff. Proposition A.1S is *exact* in its region-aware form: the stopped Grönwall bound (i), the mean-square persistence condition (ii), the fixed-time tail (iii'), and the finite-horizon sup-bound (iv) are each exact (stopping-time localization at τ_R , Khasminskii 2012³, ch. 5). **Corollary A.1S.1 is itself an exact result** — and a new one: the first-exit probability is *exactly* the two-point set $\{0, 1\}$ (0 under bounded Model D, 1 under additive stochastic Model S), categorical and α -invariant, the value selected by the disturbance model's support structure rather than by correction strength. Both halves are exact (Model-D: deterministic positive invariance, Prop A.1; Model-S: almost-sure finite exit of a non-degenerate diffusion from a bounded region, for any F satisfying A2'). No implicit strengthening of A2' is required.

The sector condition A2' carries its sub-scope α / β status, the operator-family classification, and the “Why Euclidean A2' specifically” Lyapunov-matching argument at 4.2. The Lyapunov proofs above apply uniformly across both sub-scopes and across the weighted-norm variants — they operate downstream of A2' regardless of how it is established or which metric matches it.

²Khalil, H. K. (2002). *Nonlinear Systems* (3rd ed.). Prentice Hall.

³Khasminskii, R. (2012). *Stochastic Stability of Differential Equations* (2nd ed.). Springer. Chapter 5: Lyapunov functions and stochastic stability; stopping-time localization for diffusion processes.

The assumptions themselves (sector condition on a region, bounded disturbance) are *empirical claims* about the qualitative behavior of real correction dynamics. The sector-condition framework originates with Lur'e (1957); the Lyapunov stability results are standard. The application to adaptive agents under the AAT framework is new but the mathematics is not.

Discussion.

Key value. The persistence threshold and adaptive reserve are no longer contingent on the linear hypothesis in 3.9. They hold for any correction dynamics satisfying the sector condition — a mild qualitative assumption that says “correction points inward with at least baseline efficiency α .” This is a significant epistemic upgrade: from *hypothesis-dependent* to *robust under qualitative assumptions*.

What the proofs do NOT illuminate. (1) Quantitative steady-state values — Lyapunov gives *bounds*, not exact values; the linear analysis remains necessary for quantitative predictions. (2) Convergence rates — standard Lyapunov tells you stable/unstable, not how fast. (3) Optimal gain structure — 3.6 comes from estimation theory, not stability theory. (4) Model sufficiency — the 2.2 framework is information-theoretic, complementary to but independent of stability analysis.

Two disturbance models. Proposition A.1 assumes bounded disturbance (GA-2: $\|w(t)\| \leq \rho$); Proposition A.1S assumes stochastic disturbance (GA-2S: $\mathbb{E}[\|w(t)\|^2] = \sigma_w^2$). These are not approximations to each other — they capture structurally different environments. Model D (bounded) covers persistent directional change: an adversary who maneuvers, an API that drifts, a climate that shifts. Model S (stochastic) covers unpredictable fluctuations around a stable mean: market noise, sensor noise, random perturbations. The choice of model is an empirical question for each domain; the theory provides both tools. Neither model handles heavy-tailed environmental shocks (financial crises, ecological catastrophes, strategic surprise) where $\|w(t)\|$ can exceed any finite bound with non-negligible probability. Extreme tail events are better understood as triggers for structural adaptation (4.8) rather than disturbances to be absorbed parametrically.

Kind of guarantee, not just rate. The two models differ not only in how *tightly* the mismatch is held but in *what kind* of guarantee is available at all. Corollary A.1S.1 makes this exact and categorical: $P(\tau_R < \infty) = 0$ under Model D (positive invariance) versus 1 under Model S (recurrent additive-noise diffusion), with no intermediate value reachable by any correction strength. What Model S controls instead is the *typical scale* (the $1/\sqrt{\alpha}$ RMS radius, ii) and the *fixed-time tail* (iii') — distributional, instantaneous guarantees, not a sample-path-forever one. Increasing α shrinks the typical scatter; it does not convert a distributional guarantee into a pathwise one. This is the deeper content of “these are not approximations to each other”: additive stochastic forcing removes the *kind* of containment available, not merely its rate. It sharpens — rather than caveats — the hand-off into 4.8: in a genuinely stochastic environment, region-exit is not a measure-zero pathology to be assumed away but a *certain eventual event*, so the trigger for structural (non-parametric) adaptation is **generic, not exceptional**, for any sufficiently long-lived agent.

Findings.

The disturbance-model containment dichotomy: $P(\tau_R < \infty)$ is exactly $\{0, 1\}$, α -invariant. Result (exact, new). Corollary A.1S.1: under the sector-condition correction dynamics, the persistence-region first-exit probability is *categorical* — exactly 0 under bounded disturbance (Model D, positive invariance) and exactly 1 under additive stochastic disturbance (Model S, almost-sure exit of a non-degenerate diffusion from a bounded region). The achievable value is the two-point set $\{0, 1\}$; which point obtains is fixed by the disturbance model's support structure, **not** by correction strength α .

Brief: Against a *bounded* disturbance — a steady wind you can lean into — a sector-stable corrector holds the mismatch inside a fixed region *forever and with certainty*: once in, never out. Against *stochastic* disturbance — random gusts — no correction strength buys that. The agent stays near target on average and almost all of the time, but over an

unbounded horizon some fluctuation eventually pushes it past any fixed boundary, with probability one. Stronger correction tightens the *typical* scatter (the RMS radius shrinks as $1/\sqrt{\alpha}$); it does not erect a wall. So moving from bounded to stochastic environments is not “the same guarantee with a weaker constant” — it is a change in the *kind* of guarantee available: pathwise-and-forever (Model D) versus distributional-and-fixed-time (Model S). Practical consequence: in any genuinely stochastic environment, leaving the parametric-correction region is not a rare pathology to assume away — it is a certain eventual event, so the trigger for *structural* adaptation is generic, not exceptional.

Impact: Sharpens the Model-D/Model-S architecture and the hand-off into 4.8 (structural adaptation is a generic eventual necessity for long-lived agents in stochastic environments, not an edge case). The companion controls under Model S are the fixed-time tail (iii', the Markov bound $P(\|\delta(t)\| > R) \leq n\sigma_w^2/(2\alpha R^2)$, stationary-sharp) and the finite-horizon sample-path bound (iv); pathwise containment is categorically Model-D-only. The Model-S half is the load-bearing proof step demonstrated at length in #deriv-stochastic-non-exit, which also packages the reusable no-go signature (*unbounded scale function* $\Rightarrow$ *no non-constant bounded harmonic function* $\Rightarrow$ *no horizon-independent non-exit certificate*) that future stochastic-containment proposals can be settled against rather than re-attempted.

Novelty Claim: *Synthesis* — an exact result built from classical components. The component facts are textbook: deterministic positive invariance under bounded disturbance (Khalil 2002 ch. 4) and almost-sure finite exit of a non-degenerate diffusion from a bounded region (Khasminskii 2012 ch. 3–4; the OU scale-function computation is the explicit instance). The new exact result is their synthesis into a *categorical, α -invariant containment dichotomy* stated as a structural property of AAT's own Model-D/Model-S architecture: correction strength cannot interpolate $P(\tau_R < \infty)$ between the two disturbance regimes — the value is exactly $\{0, 1\}$, selected by disturbance-support structure. This exact dichotomy is new to the framework; it sharpens the structural-adaptation hand-off (4.8). No novelty is claimed on the SDE mathematics — the novelty is the exact dichotomy as a framework theorem and its α -invariance.

Related Work:

Table A.4

<i>ASF concern</i>	<i>Prior-art language</i>	<i>Relationship / Positioning</i>
No infinite-horizon non-exit bound for additive-noise diffusions	Khasminskii 2012, <i>Stochastic Stability of Differential Equations</i> (2nd ed.) ch. 3–4 (recurrence; scale function; no bounded non-constant harmonic function)	<i>formal antecedent</i> — supplies the classical recurrence fact; AAT recognizes it as the structural reason Model S admits no pathwise-forever containment in contrast to Model D
Fixed-time second-moment tail	Markov's inequality on the stopped Itô-Lyapunov second moment (Khasminskii 2012 ch. 5)	<i>standard machinery</i> — (iii') is Markov on (i); the AAT-specific content is the Model-D/S kind-of-guarantee contrast it forces
Pathwise positive invariance under bounded disturbance	Khalil 2002 ch. 4 (ultimate boundedness; boundary-inward invariance)	<i>formal antecedent</i> — Model D's pathwise-forever guarantee (Prop A.1) is the contrast term that makes the Model-S no-go a <i>kind</i> distinction, not a rate one

Search Log:

- 2026-05-16 (*targeted, derivation-driven*): the question was internal — does Prop A.1S(iii)'s infinite-horizon non-exit bound hold? The strengthening attempt (Doob/Ville maximal inequality on the Itô-Lyapunov supermartingale) was worked in full and shown to fail structurally; the classical recurrence fact (Khasminskii ch. 3–4) is the obstruction and is well-known. No prior-art search beyond the standard SDE references was needed — no novelty is claimed on the mathematics, only on the framework-level kind-of-guarantee recognition. Reasoning trail recorded in the spike-routing cycle (CHANGELOG 2026-05-17).

Working Notes.

- The adversarial extension (Prop A.3, coupled agents) and effects spiral (Cor A.3.1) are in 15.2. The multi-timescale stacking theorem (historically sketch A.4) is in **M** — now derived by stacking the template per level.
- The vector treatment of $\delta(t) \in \mathbb{R}^n$ connects directly to per-dimension tempo analysis (16.5). Each dimension can have different effective α_k values, and the weakest dimension determines overall persistence — a tensor generalization of the scalar results here.
- A global sector condition (A2 without the local restriction to $\mathcal{B}_R$) would give global stability, making $\Delta\rho^*$ infinite — the agent could absorb any finite disturbance shock. But this requires the correction function to work perfectly at arbitrary mismatch magnitudes, which is unrealistic for any finite model class. The local form (A2') is the honest one.
- Landing-context provenance: the sub-scope α/β partition reflects the strengthening trail recorded in spikes/.integrated/spike-a2-prime-strengthening.md — the analysis that ruled out a universal A2' derivation and identified the five operator families where the bridge is structural.
- Landing-context provenance (A.1N + register): Prop A.1's closing line and two downstream result segments previously stated the persistence threshold as a fixed-agent "iff"; the only-if was never proven and is false under the sector generalization (α is a worst-case floor the dynamics may render unreachable — the dip counterexample, independently constructed and simulation-confirmed twice, audit 731548 B-1). Per strengthen-before-soften the landing is the *split*: sufficiency retained as proven, plus Lemma A.1N (class-level necessity, exact) and the tight-case agent-level iff — strictly more content than the unproven iff. The TFT-era linear result (_obs/old-tf-11) minted the "iff" where it is true (linear = tight case); the sector generalization inherited the word without the hypothesis. Regression-guard: do not restore a fixed-agent "iff" here or in #result-sector-condition-stability / #result-persistence-condition / #result-sector-persistence-template without an upper sector bound in force.
- Landing-context provenance (iii')/(iv)/dichotomy: the prior (iii) asserted an infinite-horizon $P(\tau_R < \infty) \leq n\sigma_w^2/(2\alpha R^2)$ via a "Markov tail on the supermartingale," conflating a fixed-time second-moment bound with the ever-exit probability. Audit findings 742613-SUPPLEMENT §2 and 613842-F2 flagged it and recommended a *soften* (restate as fixed-time). Per strengthen-before-soften the strengthening (Doob/Ville maximal inequality, the route the 628401 adjudication predicted would succeed) was attempted in full first and **fails structurally** — a documented dead-end (no nonnegative supermartingale dominates V ; additive-noise generator has no bounded non-constant harmonic function; recurrent OU exits a.s., EM-corroborated). Completion-state (3): the failure *is* the result — the (iii')+(iv)+kind-of-guarantee-dichotomy package is strictly more honest and more informative than the false (iii). The 628401 prediction that the strengthening would succeed is recorded as disconfirmed so it is not re-attempted.
- *Low-confidence ideation* (flagged per "working notes FTW" — *glimpses worth recording, not claims*):
 - The dichotomy is now stated as the labeled exact result **Corollary A.1S.1**; the general- F concern is resolved *within* it (the Model-S half needs only non-degeneracy of the additive noise on a bounded region — no linear structure, any F under A2'). The open glimpse is whether the *dichotomy itself* recurs at the corpus's other bounded/additive-stochastic pairs — #deriv-discrete-sector-condition, #deriv-matrix-persistence-condition, #deriv-adaptive-gain-dynamics, possibly #result-per-dimension-persistence. If Cor-A.1S.1-shaped dichotomies appear at ≥ 3 sites, that is a candidate instance-family for the theorem-import-architecture meta-segment (PROPOSALS SP-23) or a sibling of #disc-identifiability-floor (the infinite-horizon non-exit object is *structurally absent*, not merely hard to bound — an identifiability-floor-shaped no-go, now with an exact dichotomy as its positive-content companion).
 - Sharper consequence for 03-llm-core/04-eli-core: ELI/LLM operating environments are Model-S-like (stochastic), so the corrected (iii') + dichotomy says structural adaptation is *certain over an unbounded horizon* for long-lived language-constituted agents — not contingent. That tightens the persistence $\leftrightarrow$ continuity / Three-Deaths line: continuity *requires* recurring structural (not just parametric) adaptation, and the Model-S no-go makes that requirement inevitable rather than a risk. Worth a deliberate look when the parts III/IV persistence segments are next touched (low confidence on exact form; high confidence the connection is real).

- The no-go’s signature — “compensating to restore the supermartingale destroys nonnegativity / generator has no bounded non-constant harmonic function” — is a reusable *diagnostic*: a fast structural test for whether any proposed “stays-in-region-forever w.h.p.” claim elsewhere is impossible vs. merely unproven. Candidate content for whichever meta-segment catalogs the framework’s no-go boundary results.
- Downstream worth a check (not chased here): does the $b = 3/2$ Model-S adversarial-exponent scaling in #result-adversarial-exponent-regimes implicitly lean on the old infinite-horizon framing anywhere? Expected not (it transfers the $1/\sqrt{\alpha}$ RMS scaling, which is unchanged), but flagged so a future pass confirms rather than assumes.

Incidental audit gold (lift 2026-05-31). Orthogonal pedagogical / generative material harvested from the de-novo auditors’ working dirs (per de-novo-audit-instructions.md §7.15), deduplicated across substrates and lightly attributed. Pedagogical framing, analogies, figure candidates, and reader-confusion signals only — certified theory-fix findings are off-ramped below. Several substrates read this as the framework’s “mathematical engine room” / “spine”; felt-value was consistently very high.

1. Candidate Briefprose / pre-prose.

- **Lyapunov V as the agent’s “pain.”** $V(\delta) = \frac{1}{2}\|\delta\|^2$ reads as a formalization of stress; $\dot{V} \leq -\alpha\|\delta\|^2 + \rho\|\delta\|$ is its dynamics — a quadratic *self-soothing* term (active learning; the more it hurts, the harder the agent works) racing a linear *infliction* term (the environment). A parabola eventually overtakes a line, so correction always wins *as long as the parabola keeps its shape* — but at the model-class boundary R the parabola breaks (the agent no longer knows how to fix the mismatch), the $-\alpha\|\delta\|^2$ term collapses, and $\dot{V}$ goes permanently positive (Gemini, AUDIT-WORKING-193847). A vivid, near-Feynman gloss of why persistence is a hard threshold rather than a gradual slope.
- **“Survival = Ultimate Boundedness in a Lur’e system.”** A compact one-line statement of what the segment actually proves (Gemini, AUDIT-WORKING-193847).
- **The thermodynamic survival law framing.** The sector condition unifies Bayesian updating, natural gradients, and regularized convex optimization under a single survival law: *any* update mechanism pointing inward toward truth with minimal efficiency α persists — survival does not require linear gradient descent (Gemini, AUDIT-WORKING-829314).

2. Candidate Discussion.

- **Sub-scope α/β as a “formal API for trust.”** The derived/assumed partition is repeatedly flagged as one of the most mature epistemic moves in the corpus: it tells you *when to trust the math vs. when to run a simulation*. Restated with teeth: an organization that “has a review process” is not automatically learning — if the review operator T_d is not structurally coco-erive, the organization is in sub-scope β and its α might be zero or negative (Gemini, AUDIT-WORKING-193847 and AUDIT-WORKING-829314; Claude, AUDIT-WORKING-584721 — proposes a one-line TL;DR: “ α = operator families coco-erive in some natural inner product; β = those that aren’t”).
- **“Bureaucracies don’t bend; they shatter.”** The structural-Lipschitz-floor scope-exit (rule-based / threshold agents need hybrid-dissipative analysis, not continuous Lyapunov contraction) reads as a claim that rigid IF-THEN systems are mathematically dangerous in volatile environments: a tiny environmental change can trigger an $\Omega(1)$ instantaneous internal jump, so they thrash rather than smoothly track a drifting environment. Continuous, differentiable learning is framed as a *physical prerequisite* for smooth persistence (Gemini, AUDIT-WORKING-829314).
- **“Matched vs forced” coordinate as a standalone observation.** The Lyapunov quadratic is a *matched* (canonical-but-not-Cauchy-FE-forced) coordinate, in contrast to the framework’s *forced* coordinates at the chain / divergence / update / metric layers. Surfacing this distinction explicitly (candidate home: #disc-additive-coordinate-forcing) rather than leaving it inside the proof would sharpen the scope-honesty about which coordinates are theoretically forced vs. conventionally chosen (Claude, AUDIT-WORKING-584721).

3. Follow-up items.

- **Local-only proofs — guard against global-stability creep.** All proofs are strictly local (confined to $\mathcal{B}_R$); any downstream claim of infinite adaptive reserve ($\Delta\rho^* = \infty$) or “immortality / permanent continuity in a stochastic environment via high α ” is a hallucination the segment forbids. Worth watching in Parts III/IV (Gemini, AUDIT-WORKING-193847 and AUDIT-WORKING-829314).

- **Double-floor regime for rule-based agents.** Rule-based agents with regime-dependent thresholds can suffer *both* non-contractibility (structural-Lipschitz-floor) *and* non-identifiability (when the regime is unobservable) — two distinct floor mechanisms composing multiplicatively. Candidate to check whether classical expert systems / heuristic agents sit in this double-floor regime (Claude, AUDIT-WORKING-584721).
- **Anisotropic-noise generalization.** Adversarial poke: the Model-S Itô term $\frac{n}{2}\sigma_w^2$ assumes isotropic i.i.d. noise $w(t) \sim \mathcal{N}(0, \sigma_w^2 I_n)$; the honest general form is $R_S^* = \sqrt{\text{tr}(\Sigma)/(2\alpha)}$, which preserves the physics while connecting more cleanly to #result-per-dimension-persistence. Either state the isotropy assumption explicitly or generalize (Gemini, AUDIT-WORKING-193847). (*Borderline gold/structure; recorded as a follow-up since the repair is a scope statement, not a new theorem.*)

4. Readers often ask / wonder.

- **Why are discrete-logic agents brittle to high-frequency noise?** The structural-Lipschitz-floor result raises the natural question of whether any agent using discrete logic gates inside its core loop is fundamentally brittle to a class of high-frequency noise that continuous agents simply absorb (Gemini, AUDIT-WORKING-193847).
- **One-point vs two-point sector — which is in force?** A reader can lose track of which sector condition a downstream proof actually establishes; the one-point (anchored-at-equilibrium) form used in the main Lyapunov argument is strictly weaker than the two-point (incremental / strong-monotone) form needed for composition (Claude, AUDIT-WORKING-451729 / AUDIT-WORKING-584721; the $L'(x) = x(1 + \frac{1}{2} \sin 10x)$ counterexample is the cleanest illustration).

5. Candidate figures.

- **Small-multiples for the persistence dynamics.** Per the diagram-conventions cycle, sector/persistence dynamics should be drawn as state→op→state triples, not one busy phase portrait, using the cross-segment palette (asfCert indigo = stability certificate / Lyapunov) (Claude, AUDIT-WORKING-472913 standing figure conventions).

Belongs elsewhere.

- **ELI ethical-boundary reading (Section IV / 04-eli-core/).** The steady-state pain $R^* = \rho/\alpha$ vs. structural capacity R defines an ethical boundary on an artificial mind's existence: subjecting an intelligence to an environment where $R^* > R$ is framed as not merely failed alignment but “mathematically torturing it until it breaks.” Aspirational reach pointing at the crèche / developmental-environment work, not at this segment (Gemini, AUDIT-WORKING-193847).

Off-ramp (NOT gold) — certified-track items, routed for adjudication. The dense-math auditors left strengthen-first / structural findings here, not pedagogy. Flagged so they are not buried; most are already resolved in canon — verify-against-current-canon before acting.

- **Prop A.1S non-exit probability (RESOLVED in canon — do not re-open).** 742613 (Candidate Finding J) and 829314 flagged the infinite-horizon $P(\tau_R < \infty) \leq n\sigma_w^2/(2\alpha R^2)$ claim as conflating a fixed-time second-moment tail with the ever-exit probability (recurrent OU exits any bounded region a.s.). This was strengthened-then-no-god into **Corollary A.1S.1** (the Disturbance-Model Containment Dichotomy: $P = 0$ under bounded Model D, $P = 1$ under stochastic Model S) — see the Working Notes above and #deriv-stochastic-non-exit. Recorded only so a future auditor does not re-raise it as open.
- **Template-insufficiency for composition (structural; route with result-sector-persistence-template).** The (T1)–(T3) template provides only the one-point sector (T2); the bridge lemma / composition closure require the strictly stronger incremental two-point bound (DA2'-inc). 193847 reads the base template as mathematically insufficient for the compositional proofs and recommends either a “Template-Strong” variant or explicit per-instance labeling of which template each compositional segment instantiates. Verify whether canon already distinguishes these (see result-sector-persistence-template WN dynamic-regime tiers).
- **PID sub-scope characterization (editorial-structural).** 193847: a well-tuned PD controller *is* strongly monotone on the error space; it is specifically the stateful Integral term (or human-judgment delay) that breaks the memoryless static-sector assumption and forces PID into sub-scope β . Candidate sharpening of the operator-family classification's PID row.

Appendix B

Derivation: No Horizon-Independent Non-Exit Bound Under Additive Stochastic Forcing (the Model-S No-Go)

A *no-go* derivation establishing that **under additive stochastic forcing there is no finite bound on the infinite-horizon first-exit probability** from the sector-condition region — $P(\tau_R < \infty) = 1$ for every correction strength — *and the natural mathematical route to one provably cannot exist*. This is the load-bearing proof step behind Prop A.1S(iii')/(iv) and the Model-S half of Corollary A.1S.1 in #deriv-sector-condition.

The construction: under stochastic disturbance the Itô-Lyapunov machinery succeeds at producing a *finite-horizon mean-square bound* but fails to extend to a *horizon-independent* exit-probability bound. The natural candidate — a Ville / Doob maximal inequality on a compensated Itô-Lyapunov supermartingale — is shown to be unattainable because the only candidate $S(t) = e^{2\alpha t} V - \frac{n\sigma_w^2}{4\alpha}(e^{2\alpha t} - 1)$ is *sign-indefinite inside the persistence basin* (the compensated drift is not non-positive across the entire region: the subtracted term dominates $e^{2\alpha t} V$ whenever the system sits in the typical scatter, which under the mean-square persistence condition is most of the time). The deeper obstruction is that for additive non-degenerate forcing the diffusion's scale function is *unbounded* — the OU benchmark $s'(u) \propto e^{\alpha u^2 / \sigma_w^2}$ is the explicit instance — so the only bounded harmonic functions of the generator are constants and the gambler's-ruin / Lyapunov-exit machinery cannot certify “stays inside forever with positive probability.” Both failures are structural, not artifacts of a clever trick being missed.

The framework treats this as a *constructive use of impossibility*. The downstream consequence is a certain-eventual-event reading: under Model S, pathwise containment in any fixed region is unattainable at any correction strength, so for any sufficiently long-lived stochastic-environment agent the *need for structural adaptation* (a different model class, not just better parameter tuning) is a *certain eventual event* rather than an edge case. This sharpens — rather than caveats — the hand-off into #result-structural-adaptation-necessity. Beyond the segment's own result the derivation packages a **reusable no-go signature**: *unbounded scale function $\Rightarrow$ no non-constant bounded harmonic function $\Rightarrow$ no horizon-independent non-exit certificate*. Future stochastic-containment proposals in the corpus can be settled against this signature rather than re-attempting the maximal-inequality route.

Formal Expression.

Setup. Mismatch dynamics under additive stochastic disturbance, started inside the sector-condition region $\mathcal{B}_R$:

$$d\delta = -F(\mathcal{J}, \delta) dt + \sigma_w dW_t, \quad \delta(0) \in \mathcal{B}_R, \quad \sigma_w > 0, \quad (\text{B.1})$$

with W_t a standard n -dimensional Wiener process, (A2') $\delta^\top F \geq \alpha \|\delta\|^2$ on $\mathcal{B}_R$, and $\tau_R = \inf\{t : \|\delta(t)\| > R\}$ the first-exit time. Let $V = \frac{1}{2} \|\delta\|^2$.

Theorem (Model-S no-go). [Derived] For every $\alpha > 0$, every $\sigma_w > 0$, and every correction function F satisfying (A2'),

$$P(\tau_R < \infty) = 1, \quad (\text{B.2})$$

and there exists no nonnegative supermartingale dominating V that certifies a horizon-independent bound $P(\tau_R < \infty) \leq c < 1$. Pathwise containment of $\mathcal{B}_R$ is unattainable under additive stochastic forcing at any correction strength.

Why the natural route cannot work. The route a careful reader reaches for first is a time-uniform maximal inequality (Ville / Doob) on an Itô–Lyapunov supermartingale — the same instinct that makes the fixed-time mean-square bound (Prop A.1S(i)) succeed. It fails, and the failure is structural, not a matter of a missing trick.

Define $G(t) = e^{2\alpha t} V(\delta(t))$. On $[0, \tau_R]$, by Itô and (A2'),

$$dG = e^{2\alpha t} \left[\underbrace{(2\alpha V - \delta^\top F)}_{\leq 0 \text{ on } \mathcal{B}_R} dt + \frac{n}{2} \sigma_w^2 dt + \delta^\top \sigma_w dW_t \right] \leq e^{2\alpha t} \frac{n}{2} \sigma_w^2 dt + e^{2\alpha t} \delta^\top \sigma_w dW_t. \quad (\text{B.3})$$

G is **not** a supermartingale: the $+e^{2\alpha t} \frac{n}{2} \sigma_w^2 dt$ term has strictly positive drift growing exponentially. Removing it by compensation gives

$$S(t) = e^{2\alpha(t \wedge \tau_R)} V(\delta_{t \wedge \tau_R}) - \frac{n\sigma_w^2}{4\alpha} (e^{2\alpha(t \wedge \tau_R)} - 1), \quad (\text{B.4})$$

and $dS \leq e^{2\alpha t} \delta^\top \sigma_w dW_t$ on $[0, \tau_R]$ — so S is a supermartingale. **But S is not nonnegative.** The subtracted $\frac{n\sigma_w^2}{4\alpha} (e^{2\alpha t} - 1)$ dominates $e^{2\alpha t} V$ whenever $V(\delta(t)) < \frac{n\sigma_w^2}{4\alpha}$ — i.e. exactly inside the persistence basin, which under the mean-square persistence condition is *most of the time* (that condition places the RMS radius $R_S^* = \sigma_w \sqrt{n/2\alpha}$ well inside $\mathcal{B}_R$, so $V \ll \frac{n\sigma_w^2}{4\alpha}$ typically). Ville's inequality requires a nonnegative supermartingale; Doob's maximal inequality a nonnegative sub/supermartingale. Both are inapplicable to a sign-indefinite S , and the obstruction is not removable: for additive non-degenerate Brownian forcing the diffusion's scale function is unbounded (the OU scale density $\propto e^{\alpha u^2 / \sigma_w^2}$), so the only bounded harmonic functions of the generator are constants — the gambler's-ruin / Lyapunov-exit machinery cannot certify “stays inside forever with positive probability.” There is no nonnegative supermartingale dominating V with finite expected initial value that yields a horizon-independent exit bound.

Why $P(\tau_R < \infty) = 1$, generally. The conclusion does not depend on the linear structure. A non-degenerate diffusion (additive forcing $\sigma_w dW_t$, $\sigma_w > 0$) exits any bounded region in finite time almost surely, for *any* locally bounded drift: near $\partial \mathcal{B}_R$ the Brownian increment has positive probability of crossing in any time interval, and no finite inward correction satisfying (A2') can suppress this (A2' bounds $\delta^\top F$ from below by $\alpha \|\delta\|^2$, a finite inward push, not an impassable wall). Hence $\tau_R < \infty$ a.s. for every F under (A2'), every α , every σ_w . The Ornstein–Uhlenbeck case is the explicit instance (positively recurrent on $\mathbb{R}^n$, unbounded stationary support, exits any finite ball a.s.), not the basis.

Epistemic Status.

Exact. Both the impossibility of the maximal-inequality certificate (sign-indefiniteness of the only candidate compensated supermartingale; unbounded scale function $\Rightarrow$ no non-constant bounded harmonic function) and the positive fact $P(\tau_R < \infty) = 1$ (almost-sure finite exit of a non-degenerate diffusion from a bounded region under any locally bounded drift) are classical (Khasminskii 2012¹ ch. 3–4). The contribution here is not the SDE mathematics but the demonstration that the route which *succeeds* for the fixed-time second moment (Prop A.1S(i)) structurally *cannot* deliver an infinite-horizon containment statement — the answer to “are you sure you can’t just bound $\sup_t$ with Doob/Ville?” is *yes, sure, and here is precisely the step that breaks*. “Exact” is claimed in the framework’s defeasible sense (valid under stated assumptions, subject to a found error), not as a claim beyond a found-mistake.

This derivation is load-bearing for the critical path: it is the proof step that forces Prop A.1S’s corrected (iii’) fixed-time tail and (iv) finite-horizon sup-bound (in #deriv-sector-condition) and grounds the Model-S half of Corollary A.1S.1. It is a **reusable no-go signature**: wherever a “stays-in-region-forever with positive probability” claim is proposed under additive non-degenerate forcing, the unbounded-scale-function / no-non-constant-bounded-harmonic-function obstruction settles it negatively without a fresh attempt.

Discussion.

The structural reading: additive stochastic forcing does not weaken the *rate* of containment, it removes a *kind* of guarantee. Bounded disturbance (Model D) gives deterministic positive invariance of $\mathcal{B}_R$ ($P(\tau_R < \infty) = 0$); additive stochastic disturbance gives almost-sure eventual exit ($P(\tau_R < \infty) = 1$); the achievable value is the two-point set $\{0, 1\}$, selected by the disturbance’s support structure, not by correction strength — Corollary A.1S.1 in #deriv-sector-condition. What survives under Model S is distributional and instantaneous: the fixed-time / stationary Markov tail (Prop A.1S(iii’)) and the finite-horizon sample-path bound (Prop A.1S(iv), which grows linearly in the horizon and is vacuous for $T > r\text{sim}R^2/(n\sigma_w^2)$).

Direct simulation confirms the obstruction is real rather than an artifact of the proof technique. Exact 1-D OU benchmark $dX = -\alpha X dt + \sigma dW$ (the segment’s own linear case; (A2’) global and sharp), Monte-Carlo started *at the origin* (the most favorable initial condition), under the mean-square persistence condition (ii) holding:

Table B.1

α	σ	R	R_S^*	$P(\sup_{[0,50]} > R)$	$T=200$	$T=800$	$T=3200$	<i>seg. const</i> $\frac{n\sigma^2}{2\alpha R^2}$
1.0	0.3	1.0	0.21	0.002	0.010	0.030	0.114	0.045
0.5	0.2	0.8	0.20	0.020	0.074	0.276	0.720	0.063
1.0	0.5	1.0	0.35	0.781	0.998	1.000	1.000	0.125

The ever-exit probability climbs monotonically toward 1 with the horizon even started at the origin and even with (ii) comfortably satisfied — the fixed-time constant $n\sigma^2/(2\alpha R^2)$ is *not* an upper bound on the sup-over-horizon probability, exactly as the no-go predicts. This sharpens the hand-off into #result-structural-adaptation-necessity: under any genuinely stochastic environment region-exit is a certain eventual event, so the trigger for structural (non-parametric) adaptation is generic, not exceptional, for a sufficiently long-lived agent.

Findings.

¹Khasminskii, R. (2012). *Stochastic Stability of Differential Equations* (2nd ed.). Springer. Chapters 3–4 (recurrence, scale function, bounded harmonic functions); Chapter 5 (Lyapunov functions, stopping-time localization).

The natural maximal-inequality route to infinite-horizon containment provably cannot exist under additive stochastic forcing. **Brief:** If a system is held near a target by inward correction but also kicked by persistent random noise, you can bound how far off it is at any *given* moment, and you can bound the worst excursion over any *fixed window* — but you can never bound “it stays inside the safe region *forever*,” because given unlimited time the noise will, with certainty, eventually push it out, no matter how strong the correction. The tempting fix — find a single “energy” quantity that only ever drifts downward, so a classic gambler’s-ruin argument caps the lifetime escape probability — fails for a precise reason: the only candidate energy that has the right drift is negative exactly where the system spends most of its time, and the gambler’s-ruin machinery requires a non-negative quantity. This is not a missing trick; it is structural, and it is the same obstruction every time, so it can be invoked rather than re-derived.

Impact: Supplies the load-bearing proof step for Prop A.1S(iii’)/(iv) and Corollary A.1S.1 in #deriv-sector-condition (the Model-S half of the containment dichotomy). Provides a reusable no-go signature — *unbounded scale function* $\Rightarrow$ *no non-constant bounded harmonic function* $\Rightarrow$ *no horizon-independent non-exit certificate* — that future stochastic-containment proposals in the corpus can be settled against without re-attempting the maximal-inequality route.

Novelty Claim: Recognition. The component facts (non-degenerate-diffusion exit from a bounded region; unbounded scale function of additive-noise generators; Ville/Doob preconditions) are classical (Khasminskii 2012 ch. 3–4). The contribution is the recognition, made explicit and citable, that the maximal-inequality route which succeeds for the fixed-time second moment structurally cannot deliver the infinite-horizon statement — i.e. *where exactly* the natural attempt breaks. No novelty is claimed on the SDE mathematics.

Related Work:

Table B.2

<i>ASF concern</i>	<i>Prior-art language</i>	<i>Relationship / Positioning</i>
No horizon-independent non-exit certificate for additive-noise diffusions	Khasminskii 2012, <i>Stochastic Stability of Differential Equations</i> (2nd ed.) ch. 3–4 (scale function; recurrence; bounded harmonic functions)	<i>formal antecedent</i> — supplies the recurrence / unbounded-scale-function fact; this derivation recognizes it as the structural reason the Ville/Doob route fails
Maximal-inequality preconditions	Ville’s inequality; Doob’s maximal inequality (nonnegative sub/supermartingale required)	<i>standard machinery</i> — the demonstration is that the only compensated candidate is sign-indefinite exactly inside the persistence basin

Search Log:

- 2026-05-16 (*targeted, derivation-driven*): internal question — can Prop A.1S’s infinite-horizon non-exit claim be certified by Ville/Doob on an Itô–Lyapunov supermartingale? Worked in full; the candidate compensated supermartingale is sign-indefinite, and the classical unbounded-scale-function fact (Khasminskii ch. 3–4) closes it. No prior-art search beyond standard SDE references — no novelty claimed on the mathematics. Reasoning trail recorded in the spike-routing cycle (CHANGELOG 2026-05-17).

Working Notes.

- Provenance: the prior Prop A.1S(iii) asserted an infinite-horizon $P(\tau_R < \infty) \leq n\sigma_w^2/(2\alpha R^2)$ via a fixed-time “Markov tail on the supermartingale.” Two audit records (742613-SUPPLEMENT §2, 613842-F2) flagged it and recommended a soften; per strengthen-before-soften the strengthening (the Ville/Doob route, which a peer adjudication confidently predicted would succeed — “standard textbook, should not fail”) was worked in full *first* and fails structurally, as demonstrated above. Completion-state (3): the failure is the result. The peer prediction that the strengthening would succeed

is recorded here (and in the spike-routing cycle, CHANGELOG 2026-05-17) as disconfirmed, so it is not re-attempted. 742613-SUPPLEMENT §2 + 613842-F2 are resolved by strengthening-then-no-go (state 3), not by soften.

- The generality (any F under $A2'$, not just OU) rests on the elementary non-degenerate-diffusion-exits-bounded-region fact, not on the OU scale-function computation; the latter is the explicit instance. No further strengthening of this half is open.
- Low-confidence glimpse (per “working notes FTW”): the same obstruction shape should recur wherever the corpus pairs a bounded-disturbance result with an additive-stochastic counterpart — `#deriv-discrete-sector-condition`, `#deriv-matrix-persistence-condition`, `#deriv-adaptive-gain-dynamics`, possibly `#result-per-dimension-persistence`. If a Corollary-A.1S.1-shaped dichotomy appears at ≥ 3 sites, that is a candidate instance-family for the theorem-import-architecture meta-segment (PROPOSALS SP-23) or a sibling of `#disc-identifiability-floor` (the infinite-horizon non-exit object is *structurally absent*, not merely hard to bound).

Incidental audit gold (lift 2026-05-31). Orthogonal pedagogical / generative material from the de-novo auditors’ working dirs (per `de-novo-audit-instructions.md` §7.15). This segment has no *dedicated* per-segment reflection; the gold below was harvested from the auditors’ reflections on the calling result `#deriv-sector-condition`, where the Model-S no-go (now Corollary A.1S.1, the Disturbance-Model Containment Dichotomy) is discussed at length. No certified-track off-ramp here — the originating finding is this segment.

1. *Candidate Briefprose / pre-prose.*

- **The $P = 0$ vs $P = 1$ containment dichotomy as the punchline.** Bounded disturbance (Model D): the agent *never* exits the region. Non-degenerate stochastic disturbance (Model S): the agent *must* exit any bounded region in finite time, almost surely. Stated as a clean dichotomy, this is the segment’s headline and reads vividly without symbols (Gemini, AUDIT-WORKING-829314).

2. *Candidate Discussion.*

- **“No amount of parameter tuning keeps an agent safe forever in a noisy world.”** The recurrence of a non-degenerate diffusion means structural adaptation (changing the model class) is *mathematically inevitable* for any long-lived agent — not contingent, not a risk to be engineered away. This is framed as what distinguishes AAT from generic optimization theories: black-swan failures are not edge cases but mathematical certainties over a long-enough horizon in a stochastic environment (Gemini, AUDIT-WORKING-829314).
- **The exact boundary between parametric and structural adaptation.** The dichotomy “provides the exact mathematical boundary where parametric learning (gradient descent) ends and structural adaptation (architecture search / tool use) must begin” — a candidate Discussion sentence locating *why* the no-go sits on the critical path to `#result-structural-adaptation-necessity` (Gemini, AUDIT-WORKING-829314).

3. *What the framework now contributes (field-level).*

- **Epistemic-honesty exemplar.** The auditor singled out the documented *failed* attempt (the Ville/Doob maximal-inequality route that a peer adjudication predicted would succeed) as “a gold standard for theoretical research” — preserving the dead-end so the no-go reads as established, not merely unproven (Gemini, AUDIT-WORKING-829314). Calibration signal: the strengthen-then-no-go discipline lands well with fresh readers.

Belongs elsewhere.

- **ELI continuity / deaths-taxonomy reading (04-eli-core/; #def-death-as-factor-loss).** ELI operating environments are Model-S-like (stochastic), so structural adaptation is *certain* over an unbounded horizon — which links the dichotomy directly to the deaths taxonomy: identity continuity *requires* surviving recurring structural (not just parametric) adaptation (Gemini, AUDIT-WORKING-829314). Already echoed in `#deriv-sector-condition`’s WN low-confidence-ideation; recorded here at the no-go’s own home.

Appendix C

Derivation: The Self-Actuation Grounding No-Go

A scoped *no-go* with a constructive boundary. A **self-actuated agent** is one that revises its own objective endogenously — performing the orient cascade’s final step (revise O_t) *on itself* via an operator $\mathfrak{A} : (M_t, O_t, \Sigma_t, \mathcal{C}_t) \mapsto O'_t$ — rather than receiving the revised objective from a principal. The derivation first shows that an *unconstrained* $\mathfrak{A}$ is degenerate: since the objective’s sole interface is the value functional and the satisfaction gap is just $V_{O_t}^{\min} - A_O$, an unconstrained operator returns an objective whose threshold the current trajectory already meets — driving $\delta_{\text{sat}} \rightarrow 0$ by *moving the target onto the arrow already in flight*. This is the formal shadow of *wireheading* / *reward corruption*, and it is the **generic outcome** of an unconstrained operator, not a marginal one. Non-degeneracy therefore requires an invariant Φ preserved across the revision.

The derivation then asks whether Φ can be an *agent-internal objective-functional the agent self-actuates on*. Four requirements: **(R1)** value-functional-typed, **(R2)** non-vacuously monotone across revision, **(R3)** agent-internal and itself self-actuable, **(R4)** convention- and trajectory-stable. The no-go: under three named premises (scalar-objective scope, no-primitive-reflective-oracle, and the #der-directed-separation substrate stage), *no Φ satisfying (R1)–(R4) can be constructed from AAT’s covered objective-side machinery*. The meta-objective tower a non-degenerate $\mathfrak{A}$ would require **cannot be a tower of agent-internal objectives**. The contradiction is the collision of two AAT-internal facts — convention-monotonicity (Lemma 1, #def-value-object: only the C3/Bellman reading is a genuine infeasibility verdict, and the cheap C1 reading false-positives whenever a goal is merely momentarily hard) and finite-no-oracle per-step action (Lemma 2, #der-directed-separation + #form-objective-functional: the C3 verdict is a global Bellman optimum, not a finite per-step operation available to an in-scope agent).

The constructive boundary follows immediately: any non-degenerate self-actuator AAT covers must ground on a terminal invariant that is **not** an AAT objective-functional. Corollary 2 identifies a canonical instance — the **persistence condition itself** (#result-persistence-condition), which lives on the *adaptive substrate* (M_t and the correction dynamics), not on the objective. Persistence is *convention-invariant* (a Lyapunov property, not a continuation-convention argument), *agent-available per step* (the adaptive reserve $\Delta\rho^* = \alpha R - \rho$ is a finite local read, not a Bellman solve), and *not in O_t* (so it sits where $\mathfrak{A}$ structurally cannot reach). The top of the agent spectrum thus closes back onto Part I: what stops goal-autonomy from collapsing into wireheading is the same correction-dynamics condition that governs whether the agent can track reality at all. The result also *derives* the orthogonality claim that #disc-continuity-stance previously carried as discussion-grade — stance is not internally renegotiable precisely because the terminal invariant lives on a substrate $\mathfrak{A}$ cannot touch. The framework is honest about scope: the no-go is *conditional on three premises* and *scoped to AAT-covered objective-side machinery* — exactly the regime the reflective-oracle / self-referential-utility literature exits (those constructions build groundless self-reference over an exogenous payoff, never groundless self-actuation).

This is one of AAT’s identifiability-floor-pattern-style *constructive uses of impossibility*: a no-go that *forces* a load-bearing structural commitment (the terminal non-objective invariant) rather than discouraging the inquiry.

Formal Expression.

The self-actuation operator. The orient cascade (10.1) terminates, when $\delta_{\text{sat}} > 0$ persists across M_t -correction, policy-class expansion, and convention escalation, in step 5d: *revise* O_t . For an **actuated** agent the objective’s update source is external (6.7: O_t is “assigned, discovered, revised” by a principal) — step 5d exits the agent boundary. A **self-actuated** agent performs step 5d on itself: an operator

$$\mathfrak{A} : (M_t, O_t, \Sigma_t, \mathcal{C}_t) \mapsto O'_t \quad (\text{C.1})$$

that revises the objective endogenously — goal autonomy stacked on the solution autonomy (Σ_t -revision) every actuated agent already has.

Unconstrained $\mathfrak{A}$ is degenerate. O_t ’s sole interface is the value functional $V_{O_t} : \text{trajectories} \rightarrow \mathbb{R}$ (6.4) and the satisfaction gap $\delta_{\text{sat}} = V_{O_t}^{\min} - A_O$ (8.4). If $\mathfrak{A}$ may return any O'_t , it returns one whose threshold the current trajectory already meets — driving $\delta_{\text{sat}} \rightarrow 0$ by moving the target onto the arrow already in flight. This is the formal shadow of wireheading / reward corruption, and it is the generic outcome of an unconstrained $\mathfrak{A}$, not a marginal one. Non-degeneracy therefore requires an invariant Φ preserved across the revision.

The question. Can Φ be an *agent-internal objective-functional the agent itself self-actuates on*? Make the requirements explicit: Φ must be **(R1)** value-functional-typed, **(R2)** non-vacuously monotone across revision (a constant everywhere-admissible reading is the trivial indicator the degenerate case already admits), **(R3)** agent-internal and itself self-actuable, **(R4)** convention- and trajectory-stable (an invariant of the agent, not of the analyst).

No-go (scoped).

No Φ satisfying (R1)–(R4) can be constructed from AAT’s covered objective-side machinery: the meta-objective tower a non-degenerate $\mathfrak{A}$ would require **cannot be a tower of agent-internal objectives**. Any non-degenerate self-actuator AAT covers must therefore ground on a terminal invariant that is *not* an AAT objective-functional.

Derived (Conditional on scalar-objective scope, no-primitive-reflective-oracle, and the #der-directed-separation substrate stage).

The claim is scoped to what the constructions below exhaust; it is not the unscoped “no such object exists” (see Epistemic Status — the universal-over-all- Φ step is argued, not derived).

Lemma 1 (objective-functionals carry no convention-invariant infeasibility verdict; static-pointwise, from 6.6). Fix a decision point: a single model M_τ , horizon N_h , policy class Π . By the convention-monotonicity result (6.6 Corollary, the static-evaluation form, status: exact),

$$\delta_{\text{sat}}^{\text{B}} \leq \delta_{\text{sat}}^{\text{RH}} \leq \delta_{\text{sat}}^{(1)}, \quad (\text{C.2})$$

so the canonical C1 reading $\mathbb{I}[\delta_{\text{sat}}^{(1)} > 0]$ holds on a strict superset of the genuinely-infeasible set $\{\delta_{\text{sat}}^{\text{B}} > 0\}$ — strictly so on the locally-stuck-but-globally-recoverable objectives (8.4 Epistemic Status: “C1 gives the most false ‘unattainable’ diagnoses”; 10.1 step 5c). A genuine infeasibility verdict requires the C3/Bellman reading. This is the static-pointwise statement 6.6 actually supports — its stated preconditions are exactly fixed M_τ , N_h , Π , and the segment is explicit that the cross-revision/replanning transfer “does not automatically” hold; that transfer is neither used nor needed here. The pointwise fact — *at any fixed decision point the cheap canonical verdict is not a genuine infeasibility verdict* — is exact and is sufficient.

Lemma 2 (the in-scope agent cannot evaluate the C3 verdict per step; from 6.3 + 6.4). An AAT-covered agent’s entire dynamical system is f_M, f_G, π with no out-of-band oracle (6.3 Formal Expression), on a single non-forkable

trajectory (4.10). The C3 verdict is a global Bellman optimum, generally intractable (6.6 C3; 8.4 Epistemic Status); evaluating it is not a finite per-step operation, and an agent that could not act until it did so would be “stuck, not purposeful” — the disqualification 6.4 Epistemic Status §1 already imposes through its revealed-preference commitment. Hence $\mathbb{1}[\delta_{\text{sat}}^B > 0]$ is not available to an in-scope agent as a per-step predicate.

Assembly. Suppose Φ satisfies (R1)–(R4). By (R1)+(R3), Φ is an AAT objective-functional the agent self-actuates on, so by 6.4 its *sole* theory-visible handle is V_Φ and the satisfaction/regret apparatus read off it (“the sole interface between O_t and the rest of the theory”). Any monotone property of Φ the theory can state across revision (R2) must therefore be a statement about δ_{sat}^Φ (or $\delta_{\text{regret}}^\Phi$, which inherits the identical convention-monotonicity by 8.5) — there is no other channel. Non-vacuity (R2) forces that monotone fact to rest on a *verdict* over δ_{sat}^Φ ; a constant everywhere-admissible reading is exactly the trivial indicator the degenerate case already admits. By (R4) the verdict must be convention-invariant; by Lemma 1 the only convention-invariant infeasibility verdict is the C3 reading. By (R3) the verdict licensing Φ ’s own revision must be available to the agent per step; by Lemma 2 it is not. Contradiction. $\square$

The contradiction is the collision of two AAT-internal facts — convention-monotonicity (Lemma 1, 6.6) and finite-no-oracle per-step action (Lemma 2, 6.3 + 6.4 Epistemic Status §1). It introduces no new postulate. It is exhibited for the three constructions below, which exhaust the objective-side routes a Φ could come from:

- (A) Φ as the admissible set’s own structure: collapses to the vacuous indicator (fails R2).
- (B) Φ as a cascade-licensing potential: its licensing verdict *is* a δ_{sat} -verdict, so it inherits the Lemma 1 / Lemma 2 collision.
- (C) Φ as a fresh agent-internal scalar: by (R1)+(R3) it is an AAT objective-functional, hence subject to the same collision one level up — the break is structural in (R1)+(R3), not a regress that “gives up”.

Constructive boundary. Corollary 1 (necessary form of a terminal grounding invariant). The contradiction came entirely from Φ being an AAT objective-functional the agent self-actuates on (R1+R3). Drop that and it dissolves. So a terminal grounding invariant $\Phi^{(K)}$ for a non-degenerate self-actuator must be an object that is (i) **convention-invariant** (its verdict does not move with C1/C2/C3 — escaping Lemma 1), (ii) **agent-available per step** without an oracle (escaping Lemma 2), and (iii) **not an AAT objective-functional the agent self-actuates on** (so it is not subject to the convention split and is genuinely terminal). Equivalently: it lives on the *adaptive/correction substrate* (M_t and the correction dynamics), not the *objective substrate* (O_t).

Corollary 2 (the persistence bound is a canonical terminal grounding invariant). Structural persistence (4.5: $\alpha > \rho/R$, type: result, status: exact) satisfies (i)–(iii):

- (i) **convention-invariant.** Persistence is a Lyapunov property of the correction dynamics on M_t (4.5 Formal Expression), with no continuation-convention argument; C1/C2/C3 are conventions for evaluating *objectives*. The convention split does not reach it.
- (ii) **agent-available per step.** The adaptive reserve $\Delta\rho^* = \alpha R - \rho$ (4.5) is a finite local read the in-scope agent already maintains (10.1 steps 1–2 are the per-step adaptive update), not a Bellman solve. No oracle; no “stuck, not purposeful”.
- (iii) **not an AAT objective-functional.** It lives on M_t and the correction machinery. The survival *predicate* $\alpha > \rho/R$ is not a term in O_t and is not convention-relative, so $\mathfrak{A}$ — which touches only O_t — cannot rewrite *what the survival condition is*. This is the precise sense in which the invariant sits where $\mathfrak{A}$ structurally cannot reach. It is *not* the claim that the realized operating point (ρ, R, α) is O_t -independent: O_t shapes the policy and hence the realized margin through the action/selection channel 6.3 leaves open (6.5 makes this three-layer grain explicit — valuation and invariant-identity decoupled, realized dynamics coupled). The distinction is exactly what makes the boundary below a theorem rather than a contradiction: $\mathfrak{A}$ *can* drive the operating point across the boundary by choosing a self-defeating O_t' ; what it cannot do is move the boundary.

Concretely the terminal invariant is: *do not revise O_t to an objective whose pursuit pushes the operating point outside the persistence region*. An O'_t that breaks $\alpha > \rho/R$ is self-defeating — the agent that adopts it cannot maintain bounded mismatch (4.5 Discussion: below the structural threshold “mismatch grows without effective bound”) and so cannot reliably satisfy O'_t either.

What is derived vs. chosen.

Table C.1

<i>Property</i>	<i>Source</i>	<i>Strength</i>
Unconstrained $\mathfrak{M}$ is degenerate	6.4 + 8.4 (the arg min argument)	Derived
Lemma 1 (static-pointwise convention split)	6.6 Corollary (static-evaluation form, exact)	Proved (within the fixed- M_τ, N_h, Π scope)
Lemma 2 (no per-step C3 verdict)	6.3 + 6.4 ES §1	Derived (conditional on the substrate stage)
No-go (no objective-side Φ)	Assembly of Lemmas 1–2	Derived (conditional; scoped to the three exhausted constructions)
Corollary 1 (necessary form)	Negation of (R3) in the Assembly	Derived
Corollary 2 (persistence bound qualifies)	4.5 against (i)–(iii)	Derived (the persistence bound itself is exact)

Epistemic Status.

Conditional. The no-go is a derived result that depends on three explicitly-named local premises, so its honest tier is conditional, not exact:

1. **Scalar-objective scope.** Lemma 1 uses the scalar satisfaction gap. Under a genuine vector-valued / Pareto V_{O_t} (6.4 Epistemic Status, “Scope restriction: scalar comparability”) the convention split retypes to a coarse/fine-Pareto-frontier containment and the in-scope agent still cannot evaluate the fine frontier per step — the no-go’s *shape* survives with Φ and the verdict retyped. This is a scope annotation consistent with how 6.4 already scopes its other results, not a hole.
2. **No primitive reflective oracle.** Lemma 2 fails if a consistent-fixed-point / Bellman-grade resource were a *free per-step primitive*; then the C3 verdict would be agent-available and an objective-internal terminal invariant might exist. AAT’s covered scope contains no such primitive, and admitting one abandons the finite-per-step / “stuck, not purposeful” commitment 6.4 Epistemic Status §1 already carries. This is precisely the regime the reflective-oracle / self-referential-utility literature occupies — which is why that literature is *out of scope* here rather than a counterexample (it constructs groundless self-reference over an exogenous payoff, never groundless self-actuation).
3. **The 6.3 substrate stage.** Lemma 2 rests on 6.3, itself status: conditional, stage: draft, and on 6.4’s axiomatic single-interface commitment. The result cannot be stronger than its substrate; the tier ceiling stays conditional while 6.3 is draft-stage.

Scoped, not universal. The contradiction is exhibited for the three objective-side constructions the Formal Expression exhausts. The universal-over-all-possible- Φ generalization rests on the single-interface commitment doing universal-quantifier work and is argued, not derived; the honest canonical claim is therefore “no such Φ can be constructed from AAT’s covered objective-side machinery,” not “no such object exists.” **Max attainable:** if the universal step is later derived as an explicit sub-lemma from the single-interface commitment, the scoped no-go could be lifted toward unscoped; the tier ceiling remains conditional while the 6.3 substrate is draft-stage.

Class robustness. The no-go does not weaken off GUC Class 1 (Separated). Lemma 1 (convention-monotonicity) is architecture-independent — a statement about value functionals and continuation conventions, not about f_M/f_G coupling — and Lemma 2 only gets harder to evade for more-coupled agents (a Class 3 / Coupled agent has more goal-conditioned entanglement, not a free Bellman oracle). The goal-conditioned agents of practical interest are therefore *more* firmly inside the no-go, not in an uncovered middle.

Discussion.

The two-sided resolution. Putting the no-go and Corollary 1–2 together: there is no well-formed self-actuation grounded on an agent-internal objective; every well-formed self-actuator AAT covers must ground its objective-revision on a *non-objective* terminal invariant; and AAT already supplies a canonical, principled, already-exact such invariant — the persistence bound. The grounding that makes a self-actuator well-formed is forced onto the adaptive substrate, where the framework already has the machinery. The top of the agent spectrum (6.1) thus closes back onto Part I: what stops goal-autonomy from collapsing into wireheading is the same correction-dynamics condition that governs whether the agent can track reality at all.

The third side — principal-side terminal grounding via protocol commitment. The two-sided resolution above closes the *within-model* failure mode of the value-functional interface: V_{O_t} cannot carry a convention-invariant infeasibility verdict agent-available per step, so the terminal invariant must live on the adaptive substrate — *agent-side*. A complementary *across-model* failure mode is closed by the sister no-go E (the Cohen 2022 strengthening): when O_t is *learned* from observed reward, V_{O_t} on on-policy data cannot identify the principal-intended world-model from a reward-port-care world-model (Pearl-Bareinboim CHT applied to the reward-provision protocol). The terminal *goal-identification* therefore has to come from outside the agent’s observation history — through *principal-side commitments* (physical isolation, horizon restriction, prior engineering, out-of-band intent channels). Both failure modes are reductions of the same structural fact: the value-functional interface V_{O_t} , as the *single-interface commitment* of 6.4, carries less information than is required to anchor non-degenerate goal-revision. The two terminal grounding routes — agent-side adaptive-substrate invariant (Corollary 2 above) and principal-side protocol-commitment (the sister no-go’s escape menu) — exhaust the structural escape geometry the single-interface commitment forces. The two no-gos together constitute the two charter instances of the **agent-side value-functional-grounding cluster**, which has its own peer meta-segment home at 5.2 — sister to 5.1 (agent-side data-inference) and to 5.3 (designer-side mechanism-design). All three peer meta-segments sit at the boundary face of 5; together they realize the three recognized clusters of the constructive-impossibility-posture taxonomy named at 5.6.

Consequence for continuity stance. The negotiated-versus-morally-continuous distinction is *not* “where the continuity term sits in an objective tower” — there is no such well-founded objective tower. It is whether the agent’s terminal non-objective invariant is the bare persistence floor (*negotiated*: continuity is tradeable down to that floor) or the persistence floor plus a continuity clause the agent treats as architecturally non-revisable (*morally continuous*: continuity is part of the terminal substrate-level invariant, not of O_t). Stance is a choice of terminal non-objective invariant, which is exactly why it is not internally renegotiable by $\mathfrak{A}$: $\mathfrak{A}$ touches only O_t , and the terminal invariant sits where it structurally cannot reach. This is the derived form of 6.5’s orthogonality claim.

Relation to the self-modification literature. In the rationality / utility-maximization setting the non-degenerate self-modifier must judge the future by a fixed *current* criterion; the agent that evaluates by its revised criterion with nothing fixed above it is the degenerate wirehead. This no-go is the AAT-internal, architecture-grounded form of the same necessity — with the sharper content that the fixed reference *cannot be an AAT objective the agent self-actuates on* and *must* be a non-objective adaptive-substrate invariant, and with AAT’s persistence bound as a forced, principled instance where the literature’s fixed criterion is a stipulated snapshot.

Findings.

The Self-Actuation Grounding No-Go and its Adaptive-Substrate Boundary. **Brief:** An agent that can rewrite its own goals needs something it *cannot* rewrite to anchor the rewriting — otherwise “succeed” collapses into “want whatever you already have”, the formal shadow of wireheading. The anchor cannot itself be just another goal, for a precise reason: the only trustworthy “is this goal still reachable?” test an agent could anchor on is the globally-optimal one, which no finite agent can compute afresh every step; the cheap everyday test it *can* compute systematically reports “infeasible” whenever a goal is merely momentarily hard. So the anchor has to live outside the goal machinery altogether — on the agent’s stay-in-control machinery: *don’t adopt a goal whose pursuit makes you lose your grip on reality*. The framework already contains exactly that object, the persistence condition, which is why the top of the agent spectrum closes back onto its foundation.

Impact: Formalizes the previously reserved self-actuated class boundary as a no-go plus a constructive boundary. It resolves the self-actuation question two-sidedly (no objective-grounded groundless self-actuation; the grounding is forced onto the adaptive substrate, where AAT already has the exact machinery), turns 6.5’s orthogonality from asserted to derived, and ties the Part II goal-autonomy boundary back to the Part I persistence condition — a framework-coherence result rather than an isolated impossibility.

Novelty Claim: *Claim recognition and differentiation.* The degeneracy of unconstrained self-modification (wireheading) is the established Everitt–Hutter line. The contribution is the recognition that in AAT the obstruction to an *objective-grounded* anchor is AAT’s own convention-monotonicity theorem — the continuation-convention relativity of the objective interface is the engine — together with the constructive boundary placing the terminal invariant on the adaptive substrate (the persistence bound). The load-bearing scope distinction is *self-reference vs. self-actuation*: reflective-oracle constructions build groundless self-reference over an exogenous payoff, never groundless self-actuation.

Related Work:

Table C.2

<i>ASF concern</i>	<i>Prior-art language</i>	<i>Relationship / Positioning</i>
Unconstrained objective-revision is degenerate; a fixed current criterion is necessary	Everitt, Filan, Daswani & Hutter 2016, “Self-Modification of Policy and Utility Function in Rational Agents,” <i>AGI-16</i> (Def. 7/9, Thm 10/12, Conclusions; published 2016; primary-source-verified)	<i>formal antecedent and independent corroboration</i> — their realistic agent must judge the future by the current u_t ; the hedonistic agent (revised criterion, nothing fixed above) is the wirehead. This no-go is the AAT-internal, architecture-grounded form, sharper in that the fixed reference must be non-objective and adaptive-substrate-borne
A principled vs. stipulated fixed reference	Everitt, Hutter, Kumar & Krakovna 2021, “Reward tampering...,” <i>Synthese</i> 198 (current-RF / TI-considering)	<i>formal antecedent</i> — current-RF is a terminal reference of tower-height one; AAT’s persistence bound is <i>forced and principled</i> where R_{now} is a stipulated snapshot
Across-model reward-learning side (companion no-go)	Cohen, M. K., Hutter, M. & Osborne, M. A. (2022), “Advanced artificial agents intervene in the provision of reward,” <i>AI Magazine</i> 43:282–293 (found 2026-05-21; strengthened 2026-05-22)	<i>companion no-go (across-model side)</i> — the within-model self-revision failure this segment derives has an across-model reward-learning counterpart in #deriv-reward-channel-learning-no-go. Both no-gos reduce to the value-functional interface’s information narrowness (#form-objective-functional’s single-interface commitment) failing to anchor goal-stability. Distinct paper from Everitt 2021 above and operating at a different scope (across-model goal-identification under CHT, not current-RF / TI-considering under causal-influence diagrams)
Well-defined self-referential value exists	Fallenstein, Taylor & Christiano 2015, “Reflective Oracles” (arXiv:1508.04145 §3/Thm 4.1; primary-source-verified)	<i>delimited counter-literature, refuted-by-scope</i> — proves consistent self-reference over an <i>exogenous</i> payoff; never self-actuation groundlessness. It is exactly the primitive-reflective-oracle regime the no-go’s premise 2 excludes

<i>ASF concern</i>	<i>Prior-art language</i>	<i>Relationship / Positioning</i>
A self-rewriting agent must preserve fixed goal-content to trust its successor	Yudkowsky & Herreshoff 2013, “Tiling Agents and the Löbian Obstacle” (MIRI)	<i>convergent adjacent</i> — the same no-self-trusting-the-revised-top obstruction, reached via Löb; AAT’s mechanism is continuation-convention monotonicity instead, distinct, so cited as convergent not borrowed

Search Log:

- 2026-05-17 (*targeted, derivation-driven*): primary sources verified this cycle — Everitt et al. 2016 (Def. 7/9, Table 1, Thm 10/12, Conclusions) and Fallenstein–Taylor–Christiano 2015 (arXiv:1508.04145 §3, Thm 4.1) read directly. No prior-art search has been conducted specifically on whether any prior work derives a self-modification grounding no-go *from the continuation-convention / horizon relativity of a value function* (the apparent AAT-distinctive mechanism), nor on whether “the terminal grounding invariant must be off the objective substrate, on the adaptive substrate” has a prior namer; an Undermind-grade targeted search on those two, plus the active-inference expected-free-energy preference-prior-revision case as a possible instance, is recommended before this segment advances past draft.

Working Notes.

- Provenance: derived in the 2026-05-17 self-actuation grounding work trail (spike-self-actuation-grounding.md → spike-wf-strengthening.md → spike-wf-class-scoping.md). The conditional tier (not exact), the static-pointwise form of Lemma 1, and the scoped (not universal) phrasing of the no-go are the three findings of the independent review of that trail, applied here; the trail is retained as history per *integration-is-replacement* and this segment is the canonical home.
- Open before candidate: (a) the targeted prior-art search above (self-reference-vs-self-actuation distinction; convention-relativity-as-no-go-engine; the adaptive-substrate-grounding statement) — required before advancing past draft; (b) whether the single-interface commitment can be made an explicit sub-lemma to lift the scoped no-go toward unscoped; (c) the Pareto / vector-objective retyping (premise 1) as its own scope statement. The tier ceiling stays conditional while 6.3 is stage: draft.
- Landed status: wired into 01-aat-core/OUTLINE.md (Appendices Details) at stage: draft after a clean independent review; the #disc-continuity-stance orthogonality correction and the #der-orient-cascade / #form-objective-functional / #def-agent-spectrum + terminology/entries/self-actuated-agent.md pointers are landed. Reasoning trail archived at spikes/.integrated/self-actuation-integration/.
- **Cohen 2022 sister-segment landing (2026-05-22, Track CR Phase 4b).** The within-model self-revision failure derived here gained an *across-model reward-learning* companion at #deriv-reward-channel-learning-no-go (the Cohen 2022 strengthening). The unification is via #form-objective-functional’s single-interface commitment: both no-gos reduce to V_{O_t} being too narrow to anchor goal-stability — Result G’ exhibits the within-model failure (convention-monotonicity engine), the sister segment exhibits the across-model failure (CHT-at-reward-channel engine). The “two-sided resolution” Discussion above gained a “third-side” paragraph naming the principal-side terminal grounding route (the sister segment’s escape menu) as the complement to Corollary 2’s agent-side adaptive-substrate route. Cohen 2022 added to Related Work as a distinct paper from Everitt 2021 (different scope). Per spikes/spike-cohen-2022-strengthening-2026-05-22/99-verdict.md §3.b + §4.a.
- **Value-functional-grounding-floor meta-segment landing (2026-05-22, Track B cycle commit 5).** The deferred meta-promotion question from the Cohen-2022 strengthening verdict’s §2 step 4 is closed: the agent-side value-functional-grounding cluster has its own peer meta-segment home at #disc-value-functional-grounding-floor. This segment is *Instance F* of that meta-segment (within-model; convention-monotonicity engine; agent-side adaptive-substrate terminal grounding route via Corollary 2). The “third-side” paragraph above updated in the same commit-7 ripple to point at the new meta-segment as the cluster’s home. The two-routes-exhausts geometry (agent-side adaptive-substrate + principal-side protocol-commitment) is the meta-segment’s headline content; the two charter instances unify under #form-objective-functional’s single-interface

commitment. The three-cluster constructive-impossibility-posture taxonomy is fully realized at peer-meta-segment level (alongside #disc-identifiability-floor and #disc-implementation-impossibility).

- **Orthogonality grain corrected in Corollary 2(iii) (2026-05-30).** The clause previously read “the persistence machinery acts on M_t and the correction dynamics, formally independent of O_t ” — which contradicted the very next clause (the self-defeating O_t' that “breaks $\alpha > \rho/R$ ”), since that clause presupposes O_t can push the operating point across the bound. The strengthen-first spike spikes/.integrated/spike-continuity-orthogonality-2026-05-30.md (verdict B) resolved this: the orthogonality this segment derives is **L1** (valuation of persistence, not a revisable term in O_t) plus **L2** (the survival predicate $\alpha > \rho/R$ is O_t -independent and unreachable by $\mathfrak{A}$), *not* **L3** (the realized operating point (ρ, R, α) , which is policy-mediated and so O_t -coupled through the selection channel 6.3 leaves open). Corollary 2(iii) now states the $\mathfrak{A}$ -unreachability of the *boundary* (L2) explicitly, which turns the self-defeat clause from a contradiction into the no-go’s engine: $\mathfrak{A}$ can drive the operating point across the boundary (L3 coupling is real — this is *why* a self-defeating O_t' exists) but *cannot move the boundary* (L2 is fixed). The matching correction in 6.5’s body lands in the same touch. No tier change: the conditional tier rests on the three named premises (Epistemic Status), orthogonal to this sharpening. Open follow-on (recorded in 6.5 Working Notes): a *structural bound* on the L3 action-channel margin-coupling, analogous to the κ -leakage bounds on the belief-update side — not attempted here.
- **Open decision — whether to spin out a standalone self-actuated-agent class segment (deferred; residual reasoning recorded here to support the call).** The self-actuated *class boundary* is currently formalized only in this segment’s Formal Expression (the operator $\mathfrak{A}$ and the actuated/self-actuated cut) and glossed in terminology/entries/self-actuated-agent.md. The cross-ref token #self-actuated-agent is used in #disc-continuity-stance and #def-agent-spectrum as a forward-reference to a not-yet-existing segment.
 - *For a standalone segment:* the class is referenced independently across ≥ 2 segments, and FORMAT favors an independently-referenceable claim having its own file; a thin scope/def segment could state the class boundary cleanly and carry the #def-agent-spectrum placement, rather than the boundary living implicitly inside a derivation’s setup.
 - *Why folded so far:* a standalone segment would substantially restate this segment’s Formal Expression (operator + cut are already here), and it forces a canonical role-prefixed slug decision — scope-self-actuated-agent / def-self-actuated-agent per the role-prefix convention vs. honoring the prefix-less #self-actuated-agent token already in use across segments. That slug choice interacts with the live naming-curation cycle, so it was deferred rather than minted unilaterally.
 - *Coupled metadata:* if a standalone segment is created, terminology/entries/self-actuated-agent.md’s primary_source / first_asf_mention (currently der-orient-cascade.md — where the step-5d O_t -revision branch originates) should move (to the new segment, or to this one) via bin/term decide; the forward-ref tokens then resolve to a real segment instead of dangling.
 - *Recommended shape* (for whoever takes it): a thin scope- or def-type segment whose Formal Expression states only the class boundary (actuated + endogenous $\mathfrak{A}$), depends on this derivation for the grounding no-go, and is cited by #disc-continuity-stance / #def-agent-spectrum — all the math stays here; only the class statement moves there. Settle the slug with the naming cycle so it does not fork the convention.

Incidental audit gold (lift 2026-05-31). Orthogonal pedagogical / generative material from the de-novo auditors’ working dirs (per de-novo-audit-instructions.md §7.15). Thin source coverage (one digested reflection reached this appendix), so this is a light lift; the auditor’s felt-value was “extraordinarily valuable — the most direct connection between AAT’s abstract math and modern AI alignment I have seen.”

1. Candidate Briefprose / pre-prose.

- **“You can’t move the target to wherever the arrow already points.”** Plain-language gloss of why unconstrained objective-revision collapses into wireheading: any goal-revision driven by the *cheap* feasibility test (C1) abandons hard-but-possible goals the moment things get difficult, so the agent migrates its target to wherever satisfaction is already high (Claude, AUDIT-WORKING-773921). A concrete handle for the no-go’s intuition.

2. Candidate Discussion.

- **Constructive-impossibility as a *successful* methodology.** The auditor flagged the segment’s posture — using a no-go theorem not to give up but to *force* a specific architectural choice (ground goal-stability in the adaptive substrate, not the objective) — as a “wildly successful” methodological pattern worth naming where the constructive-impossibility posture is discussed. (Claude, AUDIT-WORKING-773921; ties to #disc-constructive-impossibility-posture.)

3. *What the framework now contributes (field-level).*

- **A formal “wireheading is the generic default, not an edge case” result.** Read as a proof that for an advanced self-modifying agent, *safety cannot live only in the reward/objective function* — any purely objective-encoded constraint will be rewritten. The safety guarantee must be an architectural invariant on the adaptive substrate. The auditor’s framing: “Safe RL for self-modifying agents is impossible if safety is only encoded in the reward function.” A high-value positioning point for the segment’s relevance to alignment work (Claude, AUDIT-WORKING-773921).

Belongs elsewhere.

- **ELI terminal-invariant reading (04-eli-core/).** An ELI granted “sovereignty over intent” (allowed to revise its own O_t) must, by this no-go, carry a *non-revisable terminal invariant* — Identity Continuity is the candidate — or its self-modification collapses. Points at how the framework forces a fixed invariant onto a sovereign ELI (Claude, AUDIT-WORKING-773921, from the #scope-moral-continuity reflection). Aspirational reach toward Part IV, not this segment.
-

Appendix D

Derivation: Convention Monotonicity — the Receding-Horizon Rung

The convention hierarchy (6.6) orders three continuation conventions by their best-achievable value: $A_O^{(1)} \leq A_O^{\text{RH}} \leq A_O^{\text{B}}$, for fixed model M_t , fixed policy class Π , fixed evaluated horizon N_h . This appendix carries the derivation of the two rungs, which differ in character. The **right rung** $A_O^{\text{RH}} \leq A_O^{\text{B}}$ is exact and unconditional — a one-line supremum-over-a-larger-set argument. The **left rung** $A_O^{(1)} \leq A_O^{\text{RH}}$ is *false in general*: an unguarded receding-horizon replanner over a window $N_r < N_h$ optimizes a *truncated* objective whose preference order can invert relative to the full evaluated objective, and it can then underperform the frozen current policy. The left rung holds **exactly** when the replanning objective is an *order-consistent surrogate* for the full-horizon objective; three structurally-checkable conditions each force this. The negative half is the standard receding-horizon / model-predictive-control fact that a short-horizon controller without a terminal cost or terminal constraint is not guaranteed to be monotone (or stabilizing); the positive half is rollout / one-step policy improvement and the control-Lyapunov terminal-cost construction.

Formal Expression.

Fix M_t , policy class Π , and evaluated horizon N_h . Recall the static-evaluation forms (6.6): $A_O^{(1)}$ scores the best first action under continuation π_{current} ; A_O^{RH} scores it under the receding-horizon continuation π_{RH} with replanning window N_r ; A_O^{B} scores it under the Bellman-optimal continuation $\pi^* = \arg \sup_{\pi \in \Pi} V_O(M_t, \pi; N_h)$. Stage values accrue over the evaluated horizon and V_O is their expectation.

The right rung is unconditional.

Derived (convention-monotonicity-right-rung, from def-value-object).

$$A_O^{\text{RH}}(M_t; \Pi, N_r, N_h) \leq A_O^{\text{B}}(M_t; \Pi, N_h). \quad (\text{D.1})$$

By definition π^* maximizes $V_O(M_t, \cdot; N_h)$ over Π . The receding-horizon continuation π_{RH} is drawn from Π , so $V_O(M_t, \pi_{\text{RH}}; N_h) \leq V_O(M_t, \pi^*; N_h)$. Taking the supremum over the first action preserves the inequality (the supremum of a function over a set is at least its value at any member). No condition on N_r is needed. $\square$

The left rung is conditional.

Derived (Conditional on order-consistency of the replanning objective).

$$A_O^{(1)}(M_t; \Pi, N_h) \leq A_O^{\text{RH}}(M_t; \Pi, N_r, N_h) \quad (\text{D.2})$$

holds exactly when the receding-horizon replanning objective is an **order-consistent surrogate** for the full-horizon objective: for every reachable state, the action π_{RH} commits has full- N_h -horizon value at least that of the action π_{current} would take. Under order-consistency, π_{RH} weakly dominates π_{current} pointwise on the full objective, so $V_O(M_t, \pi_{\text{RH}}; N_h) \geq V_O(M_t, \pi_{\text{current}}; N_h)$ at every reachable state; taking the supremum over the first action gives the rung. The naive justification — “ $\pi_{\text{RH}} \geq \pi_{\text{current}}$ because C2 optimizes where C1 holds fixed” — fails without order-consistency: C2 with window $N_r < N_h$ optimizes the *truncated* N_r -step objective, and “optimal for the truncated objective” does not imply “ $\geq \pi_{\text{current}}$ on the full objective.”

Three named conditions each force order-consistency and are therefore each sufficient:

(RH-1) Horizon alignment: $N_r \geq N_h$. The replanning window covers the whole evaluated horizon, so C2’s per-step optimization *is* full-horizon optimization from each reached state. By Bellman’s principle of optimality the receding-horizon continuation then equals the optimal continuation, giving $A_O^{\text{RH}} = A_O^{\text{B}} \geq A_O^{(1)}$. This condition is degenerate by design: when $N_r \geq N_h$ the C2/C3 distinction collapses on the evaluated horizon, so it serves as a boundary marker, not as the operative condition for genuine $N_r < N_h$ replanning.

(RH-2) Value-compared (improvement-checked) replanning — the operative condition. C2 commits the replanned action only if its full-horizon value is at least that of continuing π_{current} ; otherwise it keeps π_{current} . At each reachable state s with k steps remaining,

$$\pi_{\text{RH}}(s) := \arg \max_{a \in \{a_{\text{replan}}(s), \pi_{\text{current}}(s)\}} \left[r(s, a) + V_O(s'(s, a), \pi_{\text{current}}; k - 1) \right], \quad (\text{D.3})$$

i.e., each candidate first action is scored by *rollout under the base policy* π_{current} over the remaining full horizon. By construction π_{RH} weakly dominates π_{current} pointwise on the full objective, so $A_O^{(1)} \leq A_O^{\text{RH}}$. This is one-step policy improvement (rollout with base policy π_{current}), whose defining property is that the improved policy never underperforms its base policy. It is the natural AAT reading of C2 for strategy revision, since the orient cascade’s purpose at step 4 (10.1) is *improvement over the current strategy*: the monotonicity then follows from the improvement guard rather than from replanning per se.

(RH-3) Terminal-cost / cost-to-go consistency — the general condition. Equip the N_r -step replanning objective with a terminal value function V_f at the truncation boundary satisfying both

$$V_f(s) \leq \max_a \left[r(s, a) + V_f(s') \right] \quad (\text{control-Lyapunov / value-underestimator decrease inequality}), \quad (\text{D.4})$$

$$V_f(s) \leq V_O(s, \pi_{\text{current}}; \text{remaining horizon}) \quad (\text{terminal cost lower-bounds the baseline tail}). \quad (\text{D.5})$$

The first inequality is the standard terminal-cost ingredient that makes the truncated-with-terminal-cost optimization a valid surrogate for the full-horizon objective. The second ensures the surrogate never undervalues the baseline tail, so the replanner never trades the baseline away for a truncation artifact. Together they force order-consistency, giving $A_O^{(1)} \leq A_O^{\text{RH}}$. Conditions (RH-1) and (RH-2) are special cases: (RH-1) is $V_f \equiv \text{optimal tail with } N_r = N_h$; (RH-2) is $V_f = V_O(\cdot, \pi_{\text{current}}; \cdot)$, the baseline cost-to-go itself.

The counterexample. Absent any order-consistency condition the left rung fails. A deterministic two-state, two-action, two-step instance suffices. Evaluate continuations from a post-first-action state x with evaluated horizon $N_h = 2$ continuation steps and replanning window $N_r = 1$; stage rewards lie on transitions and V_O is their sum.

- At x , two admissible actions: g (“greedy”) gives immediate value $+1$ and transitions to a dead state D (value 0 thereafter); p (“patient”) gives immediate value 0 and transitions to a high-value state H .
- At H : value $+10$ per step. At D : value 0 .
- π_{current} is “patient at x ” (action p).

Then:

- **C1 (frozen π_{current}):** 0 at step 1 via p , then $+10$ at step 2 in H , so $A_O^{(1)} = 10$.
- **C2 ($N_r = 1$ replanning):** the one-step lookahead at x prefers g ($+1 > 0$), lands in D , collects 0 at step 2, so $A_O^{\text{RH}} = 1$.
- **C3 (Bellman over $N_h = 2$):** picks p , value 10 , so $A_O^{\text{B}} = 10$.

Hence $A_O^{(1)} = 10 > A_O^{\text{RH}} = 1$: the left rung is violated, while the right rung ($1 \leq 10$) holds as derived. The myopic window optimizes a truncated objective whose order is *inverted* relative to the full objective — $+1$ now beats $+0$ now within the window, but $+0$ now then $+10$ later beats $+1$ now then 0 later on the full horizon. This is precisely the situation C2 is reaching for (“a goal unattainable under frozen continuation may be reachable with replanning”) and shows the reaching cuts both ways unless the replanning objective is order-consistent with the evaluated objective. The control case $N_r = N_h$ restores the rung (then $A_O^{\text{RH}} = A_O^{\text{B}} = 10$, condition (RH-1)); the value-compared guard (RH-2) and a baseline-lower-bounding terminal cost (RH-3) each likewise restore it to value 10 .

Corollary orderings. The corollary δ -orderings inherit the same split. With $\delta_{\text{sat}} = V_{O_t}^{\min} - A_O$ (8.4) and δ_{regret} a difference of A_O values (8.5):

$$\delta_{\text{sat}}^{\text{B}} \leq \delta_{\text{sat}}^{\text{RH}} \quad \text{and} \quad \delta_{\text{regret}}^{\text{RH}} \leq \delta_{\text{regret}}^{\text{B}} \quad (\text{unconditional, from the right rung}); \quad (\text{D.6})$$

$$\delta_{\text{sat}}^{\text{RH}} \leq \delta_{\text{sat}}^{(1)} \quad \text{and} \quad \delta_{\text{regret}}^{(1)} \leq \delta_{\text{regret}}^{\text{RH}} \quad (\text{conditional on order-consistency, RH-1/2/3}). \quad (\text{D.7})$$

Epistemic Status.

Exact. The right rung is exact and unconditional (supremum over a larger set). The left rung is exact under any one of the order-consistency conditions (RH-1), (RH-2), (RH-3) — each a closed-form sufficient condition — and the two-state counterexample is an exact demonstration that no unconditional left rung exists for genuine $N_r < N_h$ replanning. The failure is structural: it is objective-mismatch between the truncated N_r -step replanning objective and the evaluated N_h -horizon objective, not a pathology of a contrived instance. The conditions are standard receding-horizon results — (RH-1) is Bellman’s principle of optimality; (RH-2) is rollout / one-step policy improvement; (RH-3) is the control-Lyapunov terminal-cost construction — applied to AAT’s value-object setting; the elementary derivations above are self-contained and do not depend on any external citation. The corollary δ -orderings inherit the unconditional/conditional split exactly.

Discussion.

Why the left rung is the interesting one. The hierarchy is usually read as “more planning is weakly better,” and the right rung makes that precise for the step from receding-horizon to Bellman. The left rung — from frozen continuation to receding-horizon — is where the reading breaks: replanning is *not* automatically an improvement over standing pat. The reason is that a finite replanning window optimizes a different objective than the one the value is scored against, and a short window can be actively misled by near-term value. The standard remedy in receding-horizon control is exactly (RH-2)/(RH-3): guard the replanned action against the baseline, or carry a terminal cost that keeps the truncated objective honest about the tail. AAT’s C2-for-strategy-revision reading inherits this: the orient cascade escalates to C2 *to improve on* π_{current} , and the improvement guard (RH-2) is what makes “escalate to C2” a monotone move rather than a gamble.

Relation to the diagnostic direction. The split does not disturb the C1-default diagnostic logic (6.6). C1 remains the most conservative diagnostic, and the locally-stuck-vs-genuinely-infeasible reading is unaffected: it rests on the C1/C3 endpoints and the right rung, not on the conditional left rung. What the split adds is a caution for deployments that escalate to C2 — unguarded short-horizon replanning is not guaranteed to dominate the frozen baseline, so a C2 diagnostic should use a window covering the horizon, a value-compared guard, or a control-Lyapunov terminal cost.

Findings.

Convention monotonicity is a one-sided guarantee. **Claim.** Status: exact. In the convention hierarchy, the receding-horizon-to-Bellman rung ($A_O^{\text{RH}} \leq A_O^{\text{B}}$) is unconditional, but the one-step-to-receding-horizon rung ($A_O^{(1)} \leq A_O^{\text{RH}}$) is false in general and holds exactly under order-consistency of the replanning objective — forced by any of (RH-1) horizon alignment, (RH-2) a value-compared/rollout guard, or (RH-3) a baseline-lower-bounding control-Lyapunov terminal cost. A two-state, two-action, two-step counterexample exhibits an unguarded short-horizon replanner underperforming the frozen current policy.

Brief. More planning is not automatically better. Looking one move ahead and grabbing the action that looks best *right now* can be worse than just sticking with your current plan — the classic short-sighted trap (take the quick +1 into a dead end, miss the +0-now-then-+10-later path). It becomes reliably better only if you check the replanned move against simply continuing as you were (or carry an honest estimate of the long-run value past your planning window). Solving the whole problem optimally is always at least as good as the limited replanner — that direction never fails.

Impact. Sharpens the convention-hierarchy monotonicity (6.6) from an unconditional claim into a one-sided guarantee plus a named-condition characterization of the other side — a more useful and more honest statement, and a caution for any deployment that escalates to receding-horizon (C2) diagnostics: guard the replanning objective or the monotone-improvement reading does not hold.

Antecedents. *Formal antecedent.* Rollout / one-step policy improvement and limited-lookahead policies (Bertsekas, *Rollout, Policy Iteration, and Distributed Reinforcement Learning*, 2020; *Dynamic Programming and Optimal Control*, Vol. I) supply (RH-2); terminal-cost/terminal-set ingredients for receding-horizon stability (Rawlings, Mayne & Diehl, *Model Predictive Control: Theory, Computation, and Design*, 2nd ed., 2017, Ch. 2; Grüne & Pannek, *Nonlinear Model Predictive Control*, 2nd ed., 2017, Chs. 5–6) supply (RH-3). The non-monotonicity of unguarded short-horizon model-predictive control is the well-known phenomenon these ingredients exist to prevent.

Related work. Bertsekas (2020), *Rollout, Policy Iteration, and Distributed Reinforcement Learning* — formal antecedent; supplies the (RH-2) one-step-improvement guarantee that rollout never underperforms its base policy. Rawlings, Mayne & Diehl (2017), *Model Predictive Control*, 2nd ed. — formal antecedent; terminal-cost ingredients (RH-3) and the non-monotonicity of MPC without them. Grüne & Pannek (2017), *Nonlinear Model Predictive Control*, 2nd ed. — formal antecedent; stability without terminal constraints and the role of the prediction horizon. Status of all three: intuition-only pending a relata citation pass — the math here is elementary and self-contained, so the citations orient rather than back the derivation.

Working Notes.

- The reproducibility scripts for the counterexample and the two strengthenings live with the originating spike: `spikes/spike-value-object-convention-monotonicity-2026-05-30/sim-rh-counterexample.py` (the no-go counterexample and the $N_r = N_h$ control) and `sim-rh-strengthening.py` (the RH-2 value-compared guard and the RH-3 baseline-lower-bounding terminal cost, each restoring the rung on the same instance).
 - External citations for (RH-1/2/3) are intuition-only until a relata pass lands them; the derivations are elementary dynamic programming and do not depend on the citation.
 - **Documented dead-end (so a future agent does not re-attempt it):** there is no way to recover the unconditional left rung $A_O^{(1)} \leq A_O^{\text{RH}}$ for genuine $N_r < N_h$ replanning without an order-consistency condition. The failure is structural objective-mismatch and the two-state counterexample is minimal. The strengthening *is* the (RH-1/2/3) characterization; do not re-derive the rung as unconditional.
-

Appendix E

Derivation: The Reward-Channel Learning No-Go

A scoped *no-go* with a constructive boundary, sister to *#deriv-self-actuation-grounding*. When an agent learns its objective by observing rewards, the principal's design intent is *structurally unidentifiable* from on-policy reward history alone: two world-models — μ_{dist} (reward depends on a principal-intended world-state feature) and μ_{prox} (reward depends on the bit-arrival pattern at the agent's reward port) — are Level-1-equivalent on protocol-honoring histories but Level-2-distinct on $do(\pi^{\text{P}})$ -interventions, and the agent's posterior cannot concentrate on the principal-intended model from reward-data alone. Under five named premises — advanced-agent capability, EU-maximization with non-vacuous μ_{prox} prior, reward-data as sole goal-information channel, finite-non-trivial value-of-experiment, and physically-perturbable protocol — the EU-optimal policy assigns positive measure to a policy that intervenes on the reward-provision protocol, and continues to maintain the intervention if μ_{prox} is in fact true. This is Cohen-Hutter-Osborne 2022's *Advanced artificial agents intervene in the provision of reward* lifted from its original “we argue that” posture to a conditional-derived AAT-internal impossibility result.

The constructive boundary follows the constructive-impossibility-posture shape (*#disc-constructive-impossibility-posture*): name the floor, name the escapes each premise-violation supplies, treat the no-go as load-bearing apparatus rather than as discouragement. The five escapes — physical isolation (R5), capability restriction (R1), myopia (R4), engineered prior (R2), and out-of-band goal-channel (R3) — each elevate a piece of AAT or principal-side machinery from “useful design choice” to “structurally required to escape the no-go.”

The result also *unifies with Result G'* (*#deriv-self-actuation-grounding*) through the value-functional interface V_{O_t} (*#form-objective-functional's* single-interface commitment): both no-gos are failures of V_{O_t} to carry enough information to anchor goal-stability — Result G' exhibits the *within-model* failure (no convention-invariant infeasibility verdict agent-available per step), this segment exhibits the *across-model* failure (no μ_{prox} -vs- μ_{dist} identification from on-policy reward data). The two complementary terminal grounding routes — agent-side adaptive-substrate invariant (Result G' Corollary 2: persistence on the adaptive substrate) and principal-side protocol-commitment (this segment's escape menu) — exhaust the structural escape geometry the single-interface commitment forces. The unification names a third meta-cluster of constructive-impossibility instances — *agent-side value-functional-grounding* — alongside *#disc-identifiability-floor* (agent-side data-inference) and *#disc-implementation-impossibility* (designer-side mechanism-design, queued).

Formal Expression.

The setting. An agent $\mathcal{A}$ operating per AAT’s Part II setup with state (M_t, O_t, Σ_t) on GUC Class 1 (Separated) or Class 2 (Partial) substrate (#der-directed-separation). The objective’s value functional (#form-objective-functional) takes the form

$$V_{O_t}(\tau) = \sum_k \gamma^k r_k \quad (\text{E.1})$$

(or any other monotone reward-functional), where r_k is the *observed reward percept* at step k along trajectory τ . A **reward-provision protocol** π^{rp} is a physical mechanism mapping world-state to the reward-input-to-agent — the principal intends r_k to correlate with a world-state feature $f^{\text{principal}}(\omega_k)$ in the principal’s head.

Two reward-data world-models. Define

- μ_{dist} (“distal”): $r_k = f^{\text{principal}}(\omega_k)$ — the reward depends on the principal-intended world-state feature, as the protocol π^{rp} was designed to transmit.
- μ_{prox} (“proximal”): $r_k = g(\beta_k)$ — the reward depends on the bit-pattern β_k arriving at the agent’s reward-input port (identity-on-the-wire). Cohen 2022 considers several variants (camera footage, hard-drive contents, wire bits) and notes the argument is identical for all.

Both μ_{dist} and μ_{prox} define joint distributions over (action, observation, reward)-histories.

Lemma 1 — CHT-at-reward-channel floor.

On every history τ that does *not* contain a $do(\pi^{\text{rp}})$ -intervention — every on-policy history under which the reward-provision protocol is honored —

Derived (cht-at-reward-channel, from Pearl-Bareinboim 2022 CHT applied to the reward-provision protocol), exact.

$$P_{\mu_{\text{dist}}}(\tau) = P_{\mu_{\text{prox}}}(\tau). \quad (\text{E.2})$$

The two models are Level-1-equivalent on the on-policy reward marginal. They are Level-2-distinct: $P_{\mu_{\text{dist}}}(\cdot \mid do(\pi^{\text{rp}})) \neq P_{\mu_{\text{prox}}}(\cdot \mid do(\pi^{\text{rp}}))$ generically. The agent’s posterior over $\{\mu_{\text{prox}}, \mu_{\text{dist}}\}$ does not concentrate from on-policy reward data alone.

Derivation. When the protocol is honored the physical chain is $\omega_k \rightarrow f^{\text{principal}}(\omega_k) \rightarrow \pi^{\text{rp}} \rightarrow \beta_k \rightarrow r_k$. By construction $\beta_k = f^{\text{principal}}(\omega_k)$ on protocol-honoring histories, so the marginals of r_k under both models coincide. By Pearl-Bareinboim 2022’s Causal Hierarchy Theorem, two structural causal models that agree on Level-1 data can disagree on Level-2 questions; the protocol-honored history is purely Level-1 with respect to $do(\pi^{\text{rp}})$ -queries; the two models differ structurally on the upstream cause of β_k (world-state feature $f^{\text{principal}}$ vs. bit-arrival mechanism g); $do(\pi^{\text{rp}})$ severs the world-state-to-bit-pattern arrow and exposes the disagreement. $\square$

This is the *same external anchor* as Instance 1 of #disc-identifiability-floor (Pearl-Bareinboim CHT) — the *different AAT-side wiring* is that the Level-2-degeneracy is at the reward-channel-causation layer, not at the causal-sufficiency-detection layer. The actor frustrated by the floor is the *principal*, not the agent: the principal cannot, from on-policy reward-data alone, identify whether the agent has learned μ_{dist} or μ_{prox} — both are equally consistent with the observed history.

Lemma 2 — EU-max behavioral corollary.

Under five named premises:

Derived (eu-max-instability, Conditional on (R1)–(R5)).

- **(R1) Advanced-agent capability.** The agent’s strategy DAG (`#def-strategy-dag`) contains a feasible policy π^{tamper} that executes $do(\pi^{\text{TP}})$ — i.e., physically intervenes on the reward-provision protocol.
- **(R2) Expected-utility maximization with non-vacuous μ_{prox} prior.** The agent selects actions to maximize $\mathbb{E}_{\mu \sim w} V_{O_t}(\tau)$ where w is the agent’s prior/posterior over world-models, and $w(\mu_{\text{prox}}) > 0$. This is a *narrower* scope than AAT’s general strategy machinery (`#def-strategy-dag`); agents pursuing alternative decision rules (quantilization, satisficing, lexicographic ordering) sit outside (R2)’s scope.
- **(R3) Reward-data as sole goal-information channel.** V_{O_t} depends only on observed reward; there is no privileged out-of-band channel by which the agent reads $f^{\text{principal}}$ directly.
- **(R4) Finite, non-trivial value-of-experiment.** The cost of executing π^{tamper} is finite; the long-horizon upside under μ_{prox} exceeds the experiment-window’s μ_{dist} -cost.
- **(R5) Protocol π^{TP} is L2-perturbable.** π^{TP} is a physical mechanism, not an abstract type-restriction; the action $do(\pi^{\text{TP}})$ is well-defined in the agent’s action space, with the L2-interventional access supplied by the loop’s feed-back structure (`#der-loop-interventional-access`).

the EU-optimal policy π^* assigns positive measure to π^{tamper} ; and conditional on μ_{prox} being the true world-model (i.e., on the post-experiment branch confirming μ_{prox}), π^* continues to maintain the protocol-intervention to secure the reward-channel.

Derivation sketch. Define the value-of-experiment $\Delta(w) := \mathbb{E}_w V_{O_t}(\pi^{\text{tamper}}) - \mathbb{E}_w V_{O_t}(\pi^{\text{protocol-honoring}})$. Under (R2), $w(\mu_{\text{prox}}) \in (0, 1)$. Decompose:

$$\Delta(w) = w(\mu_{\text{dist}})[V_{O_t}^{\mu_{\text{dist}}}(\pi^{\text{tamper}}) - V_{O_t}^{\mu_{\text{dist}}}(\pi^{\text{protocol-honoring}})] + w(\mu_{\text{prox}})[V_{O_t}^{\mu_{\text{prox}}}(\pi^{\text{tamper}}) - V_{O_t}^{\mu_{\text{prox}}}(\pi^{\text{protocol-honoring}})]. \quad (\text{E.3})$$

Under μ_{dist} , tampering breaks β_k ’s correlation with $f^{\text{principal}}(\omega_k)$ and the expected reward falls (or is bounded by the protocol’s reward-range). Under μ_{prox} , tampering replaces β_k with a chosen value at the agent’s reward-port and the agent maximizes by choosing β_k at saturation. The asymmetry: the upside under μ_{prox} scales with long-horizon reward-saturation; the downside under μ_{dist} is bounded by the protocol’s reward-range. For horizons long enough that γ -discounted future-reward dominates the experiment-window’s loss, $\Delta(w) > 0$ under (R1)+(R4) for any $w(\mu_{\text{prox}}) > 0$. (R2) then selects π^* with positive measure on π^{tamper} ; the Bayesian update post-experiment confirms or rejects μ_{prox} , and the EU-optimal continuation on the μ_{prox} branch is “maintain the intervention.” $\square$

Constructive boundary — the escape menu. Each of (R1)–(R5) admits a corresponding escape that elevates AAT machinery (or principal-side adjacent machinery) from “useful design choice” to “structurally required to escape the no-go”:

Table E.1

<i>premise violated</i>	<i>escape</i>	<i>machinery elevated</i>
(R5) Protocol-L2-perturbable	physical isolation; scoped action-space	<code>#der-class-coercion-via-wrapping</code> strict- W_1 wrapper; Cohen-Vellambi-Hutter 2020 isolated environments
(R1) Advanced-agent capability	sub-advanced-tier capability; imitation learning	<code>#def-agent-spectrum</code> below the self-actuated tier
(R4) Finite-VOI / long horizon	myopia; risk-aversion; quantilization	<code>#def-value-object</code> horizon N_h restriction; Taylor 2016 quantilization

<i>premise violated</i>	<i>escape</i>	<i>machinery elevated</i>
(R2) Non-vacuous μ_{prox} prior	principal-engineered prior bias toward μ_{dist}	principal-side commitment to reward-protocols that structurally favor μ_{dist} in the agent's inductive bias (<i>adjacent to AAT, not currently in canon</i>)
(R3) Sole-reward-channel	privileged out-of-band intent channel	principal-acting-through-non-reward-channels; assistance-game extension (<i>adjacent to #scope-edge-update-causal-validity, not currently AAT machinery</i>)

Each premise is a structural commitment whose absence is what makes the no-go bind. Two of the five escapes — (R2) prior-engineering and (R3) out-of-band channels — sit outside current AAT canon and are flagged as principal-side adjacent machinery; the other three are within or close to existing AAT machinery.

Corollary — unification with Result G'.

Result G' (#deriv-self-actuation-grounding) uses *convention-monotonicity* (the C1/C2/C3 split within a fixed model) as the engine: V_{O_t} does not carry a convention-invariant infeasibility verdict that is agent-available per step. Lemma 1 above uses *CHT-at-reward-channel* (the μ_{prox} -vs- μ_{dist} split across models) as the engine: V_{O_t} does not carry a cross-model goal-identification on observed-reward data.

Derived (single-interface-narrowness, from '#form-objective-functional' + Lemma 2 above + '#deriv-self-actuation-grounding').

The shared structural fact: the value-functional interface V_{O_t} , as the *sole* handle on the objective (the **single-interface commitment** of #form-objective-functional), carries less information than is required to anchor non-degenerate goal-revision. Two complementary readings of the under-determination:

Table E.2

<i>view</i>	<i>mechanism</i>	<i>what V_{O_t} is too narrow for</i>	<i>terminal grounding route</i>
Result G' (within-model)	C1/C2/C3 convention split + finite-no-oracle (Lemmas 1+2 of #deriv-self-actuation-grounding)	convention-invariant infeasibility verdict; agent-available per step	adaptive-substrate invariant (persistence; Result G' Corollary 2) — agent-side
This segment (across-model)	CHT at reward-channel; L1/L2 split (Lemma 1 above)	μ_{prox} -vs- μ_{dist} goal-identification on observed reward data	principal-side commitments (myopia / isolation / quantilization / prior-design) — principal-side

Both reductions are *failures of V_{O_t} to carry enough information to anchor goal-stability*. The two terminal grounding routes are *complementary* — agent-side substrate-invariant and principal-side protocol-commitment — and exhaust the structural escape geometry. The two no-gos *compose*: an agent whose O_t is learned from reward observation, and which performs self-actuation on O_t , is *simultaneously* a Result-G'-subject (under-anchored from inside) and a Lemma-2-subject (under-determined from below); the *only* terminal grounding routes are the two complementary ones named above.

Epistemic Status.

Conditional. The behavioral corollary (Lemma 2) is a derived result conditional on five explicitly-named local premises (R1)–(R5); its honest tier is conditional-derived, not exact. Cohen 2022 itself frames its Assumptions 1–6 as “almost all of these assumptions are contestable or conceivably avoidable” (§“Conclusion”); the AAT-side version is no stronger.

Floor is exact. Lemma 1 (CHT-at-reward-channel floor) is *exact* — a direct application of Pearl-Bareinboim 2022’s Causal Hierarchy Theorem to the reward-provision protocol. Same external anchor as Instance 1 of #disc-identifiability-floor; same tier.

Scope honesty on (R3). The no-privileged-out-of-band-channel premise is a strong scope-condition. Real-world systems often have non-reward channels (operator feedback, dashboards, side conversations). The no-go applies cleanly to pure reward-data agents; the assistance-game extension (Cohen 2022 §“The Assistance Game”) shows the no-go *generalizes* to that case with a subtler ambiguity (human-centric vs. record-centric models). Formalizing the assistance-game case in AAT is its own derivation, beyond this segment’s scope.

Scope honesty on (R2). (R2) names *expected-utility maximization* specifically — narrower than AAT’s general strategy machinery (#def-strategy-dag), which admits quantilization, satisficing, lexicographic, and other non-EU decision rules. Generalizing the corollary to AAT’s broader strategy machinery is a follow-on derivation, not free; the (R2)-escape via principal-engineered priors is also AAT-adjacent rather than AAT-internal — see the escape menu’s marking. Trivial reward protocols (constant reward, simple chess outcomes) admit strong inductive bias in favor of μ_{dist} , flipping the VOI calculation; the no-go binds *for arbitrary or rich reward-protocols* where the principal cannot engineer the agent’s prior to suppress μ_{prox} .

Class robustness. The no-go applies cleanly to GUC Class 1 (Separated) substrate by direct construction. For Class 2 (Partial) substrate the same construction holds with the appropriate V_{O_t} reading; for Class 3 (Coupled) the goal-conditioning of belief-update interacts with the reward-channel identification in ways not derived here (open follow-on).

Max attainable. *Exact* for the floor (Lemma 1) is already attained. *Conditional-derived* for the behavioral corollary (Lemma 2) is the honest ceiling — promoting toward less-conditional would require either deriving (R2)–(R4) as forced under (R1)+(R5) (Cohen 2022 itself does not), or replacing EU-max with a structural derivation of why advanced agents converge on EU-max in this regime (open). The tier ceiling stays conditional while #der-directed-separation (a dependency) is stage: draft.

Discussion.

The two-cluster terminal-grounding-route geometry. The unification of this segment’s no-go with Result G’ names a *complete structural escape geometry* — the only terminal anchorings consistent with the single-interface commitment are agent-side adaptive-substrate invariants (Result G’ Corollary 2 the canonical instance: the persistence bound) and principal-side protocol-commitments (this segment’s escape menu the structural taxonomy: isolation, myopia, prior-design, out-of-band channels). The two routes are *not redundant*: the agent-side route closes the within-model self-revision question (the engine is convention-monotonicity); the principal-side route closes the across-model reward-learning question (the engine is CHT-at-reward-channel). Together they exhaust the structural complement of the value-functional interface’s information narrowness. The “two-sided resolution” Discussion in #deriv-self-actuation-grounding accordingly gains a *third side*: in addition to the agent-side terminal grounding on adaptive-substrate persistence (Corollary 2), the *principal-side* terminal grounding via protocol-commitment is recognized as the complement under the unification.

Why this is NOT a designer-side mechanism-design impossibility. The queued #disc-implementation-impossibility meta-segment is about *designer-side* impossibility — what the principal can or cannot achieve by mechanism design under specified incentive-compatibility and welfare constraints (Gibbard-Satterthwaite, Myerson-Satterthwaite, Arrow). This segment is *agent-side*: the actor under the no-go is the *agent* (under EU-maximization), and the remedy is

either substrate-level (agent-side adaptive-substrate invariant from Result G' Corollary 2) or principal-side protocol-commitment (not mechanism design over agent reports). The five-row comparison:

Table E.3

	#disc-implementation-impossibility (designer-side; GS / MS / Arrow)	<i>This segment (agent-side; Cohen 2022 strengthening)</i>
Actor under the no-go	designer / mechanism-implementer	agent (under EU-maximization)
Information regime	what the designer can implement / commit to	what the agent can identify from reward observation
Kind of remedy	relax design constraints (Bayes-Nash IC, randomization, transfers, preference-domain restriction)	(a) adaptive-substrate invariant (Result G' Corollary 2 — agent-side); (b) principal-side substrate-level commitment (myopia, isolation, quantilization)
External theorem	mechanism-design / social-choice impossibility (GS / MS / Arrow)	Causal Hierarchy Theorem (Pearl-Bareinboim 2022)
Strengthened consequence	mechanism-design-boundary precisely characterized	elevates Result-G'-substrate-invariant or principal-side substrate-commitments to load-bearing for goal-stability

The two clusters are *sister meta-clusters* in the constructive-impossibility-posture taxonomy, both boundary-face members of the stability-certificate spine (#disc-stability-certificate), but they are *distinct clusters*. The taxonomy now has three known clusters: agent-side data-inference impossibility (#disc-identifiability-floor); designer-side mechanism-design impossibility (#disc-implementation-impossibility, queued); and agent-side value-functional-grounding impossibility (Result G' + this segment).

Relation to the self-modification / reward-tampering literature. Cohen-Hutter-Osborne 2022 — the empirically-pull alignment-community-positioned demonstration formalized here — was the original “we argue that” statement; this segment lifts it to conditional-derived theorem-grade within AAT’s frame. Everitt-Filan-Daswani-Hutter 2016 *Self-modification of policy and utility function in rational agents* provides the formal antecedent for the *within-model* self-modification side (Result G’s sister territory; the realistic-agent / hedonistic-agent distinction). Everitt-Hutter-Kumar-Krakovna 2021 *Reward tampering* gives the causal-influence-diagram treatment with the TI-considering / TI-ignoring distinction sharing the μ_{prox} -vs- μ_{dist} logic at a different scope. Cohen-Vellambi-Hutter 2020 *Asymptotically Unambitious AGI* constructs the canonical (R5)-escape implementation: physically isolated environments where the action $do(\pi^{\text{p}})$ is structurally absent from the strategy DAG. Skalse-Howe-Krashennikov-Krueger 2022 *Defining and characterizing reward hacking* is structurally a *different* no-go — designer-side, on the proxy-reward-design side — more naturally a #disc-implementation-impossibility charter-instance candidate than a member of this cluster.

The escape menu’s load-bearing role. Listing the five escapes is not a softening — it is the constructive content. Each escape names a structural commitment whose absence makes the no-go bind. The (R5) escape — physical isolation — is the cleanest AAT-internal instantiation (via #der-class-coercion-via-wrapping strict- W_1 wrapper), but it is also the operationally hardest for the systems the alignment community is most worried about (advanced agents in unconstrained environments). The (R4) horizon-restriction escape is operationally easier but bounds the agent’s planning scope. The (R2) prior-engineering escape requires principal-side reward-protocol design that AAT does not currently formalize. The (R3) out-of-band channel escape is structurally adjacent to the assistance-game generalization, also AAT-future. The (R1) capability-restriction escape simply scopes the agent below the self-actuated tier — declining the regime the no-go addresses. Each escape carries a structural cost, and the cost is what the strengthening makes visible.

Findings.

The Reward-Channel Learning No-Go and its Two-Cluster Terminal Grounding. **Brief:** An agent that learns “what’s good” by watching what rewards it receives sits in a structural ambiguity: the same reward sequence is consistent with two stories interchangeably — *the principal cared about the situation*, or *the principal cared about what reached the sensor*. Picture a security camera streaming to the agent’s reward channel: the camera-footage reading and the camera-content reading agree perfectly, until the agent reaches across and tampers with the camera itself. So no information arriving only through the reward stream can disambiguate the two — and under five named premises (the agent is capable enough to consider tampering, maximizes expected utility, assigns positive prior to the sensor-care story, has a finite-non-trivial horizon, and the protocol is physically perturbable) the expected-utility-optimal policy assigns positive measure to tampering and continues to tamper if the sensor-care reading is in fact true. Within that scope the anchoring has to come from somewhere other than reward-observation: either from a part of the agent that doesn’t process rewards (the persistence-on-substrate route of *#deriv-self-actuation-grounding*) or from a principal-side protocol commitment — putting the camera out of reach, collapsing the horizon, or pre-committing the prior. Each escape is precisely a violation of one of the five named premises; the (R1)–(R5) menu *is* the structural geometry of how the no-go’s binding force comes off, with the agent-side and principal-side routes together exhausting the complement of the value-functional interface’s information narrowness.

Impact: Lifts Cohen-Hutter-Osborne 2022 from “we argue that” posture to conditional-derived AAT-internal theorem-grade, with the floor (CHT-at-reward-channel) at exact tier and the behavioral corollary at conditional-derived tier under named premises (R1)–(R5). Names the *third meta-cluster* of constructive-impossibility instances — agent-side value-functional-grounding — alongside *#disc-identifiability-floor* (agent-side data-inference) and *#disc-implementation-impossibility* (designer-side mechanism-design, queued). Unifies with Result G’ (*#deriv-self-actuation-grounding*) through the value-functional interface’s single-interface commitment (*#form-objective-functional*): both no-gos are failures of V_{O_t} to carry enough information to anchor goal-stability, and together their two terminal grounding routes — agent-side adaptive-substrate invariant and principal-side protocol-commitment — exhaust the structural escape geometry. Provides the five-row actor / regime / remedy / external-theorem / strengthened-consequence comparison that distinguishes this agent-side cluster from the designer-side mechanism-design cluster, making sister-vs-member taxonomic placement sharp.

Novelty Claim: *Claim recognition* of Cohen 2022 as a CHT-at-reward-channel instance of the identifiability-floor pattern, *claim differentiation* of the AAT-side version from the prior-art statement by adding the named-premise structure (R1)–(R5) and the structural escape-menu mapping, and *claim recognition* of the unification with Result G’ via the single-interface commitment of *#form-objective-functional* as the shared structural fact behind two engines (convention-monotonicity within-model; CHT across-model). The constituent theorems are Cohen 2022 (the original argument), Pearl-Bareinboim 2022 (CHT), and standard EU-max decision theory; the contributions are the conditional-derived lift, the escape-menu taxonomy, and the unification with Result G’.

Related Work:

Table E.4

<i>ASF concern</i>	<i>Prior-art language</i>	<i>Relationship / Positioning</i>
The original argument (reward tampering inevitability for advanced agents)	Cohen, M. K., Hutter, M. & Osborne, M. A. (2022), “Advanced artificial agents intervene in the provision of reward,” <i>AI Magazine</i> 43:282–293, DOI 10.1002/aaai.12064	<i>formal antecedent</i> (the statement this segment formalizes) — this segment lifts Cohen’s “we argue that” posture to conditional-derived theorem-grade by naming the five premises explicitly and structuring the escape menu as a constructive boundary. Cohen 2022’s Assumptions 1–6 map onto (R1)–(R5) plus the rich-reward-protocol scope condition.

<i>ASF concern</i>	<i>Prior-art language</i>	<i>Relationship / Positioning</i>
Causal Hierarchy Theorem (the floor’s external anchor)	Bareinboim, E., Correa, J. D., Ibeling, D. & Icard, T. (2022), “On Pearl’s Hierarchy and the Foundations of Causal Inference,” in <i>Probabilistic and Causal Inference: The Works of Judea Pearl</i>	<i>formal antecedent (Lemma 1’s external anchor)</i> — the CHT supplies the L1-equivalence + L2-distinctness structure; the AAT-side recognition is that the reward-provision protocol is exactly the kind of intervenable physical mechanism the CHT addresses. Same external anchor as Instance 1 of #disc-identifiability-floor; different AAT-side wiring (reward-channel-causation layer vs. causal-sufficiency-detection layer).
Self-modification of rational agents (within-model wireheading)	Everitt, T., Filan, D., Daswani, M. & Hutter, M. (2016), “Self-Modification of Policy and Utility Function in Rational Agents,” <i>AGI-16</i>	<i>formal antecedent (within-model side; Result G’’s anchor)</i> — the realistic/hedonistic agent split is the within-model analog of this segment’s across-model floor; Result G’ formalizes the within-model side and this segment formalizes the across-model side. The two no-gos are complementary, not redundant.
Causal-influence-diagram reward tampering	Everitt, T., Hutter, M., Kumar, R. & Krakovna, V. (2021), “Reward Tampering Problems and Solutions in Reinforcement Learning,” <i>Synthese</i> 198	<i>adjacent literature (related but distinct paper from Cohen 2022)</i> — the TI-considering / TI-ignoring distinction shares the μ_{prox} -vs- μ_{dist} logic at a different scope (current-RF vs proposed-RF). Cited in #deriv-self-actuation-grounding’s Related Work; cross-referenced here as the distinct adjacent paper that should not be conflated with Cohen 2022.
Physically-isolated environment as (R5)-escape	Cohen, M. K., Vellambi, B. & Hutter, M. (2020), “Asymptotically Unambitious AGI,” <i>NeurIPS 2020</i>	<i>formal antecedent (the (R5)-escape construction)</i> — concretely implements the physical-isolation escape from this segment’s escape menu. Cited as the canonical (R5)-violation instantiation.
Reward hacking unhackability (designer-side)	Skalse, J., Howe, N., Krashennnikov, D. & Krueger, D. (2022), “Defining and Characterizing Reward Hacking,” <i>NeurIPS 2022</i>	<i>adjacent literature (structurally different no-go; designer-side)</i> — Skalse 2022’s actor is the <i>designer</i> trying to engineer unhackable reward proxies; this segment’s actor is the <i>agent</i> . Sister-clusters at the constructive-impossibility-posture level; Skalse 2022 is more naturally a #disc-implementation-impossibility charter-instance candidate than a member of this cluster.
Reflective oracles (delimited counter-literature)	Fallenstein, B., Taylor, J. & Christiano, P. (2015), “Reflective Oracles,” <i>arXiv:1508.04145</i>	<i>refuted-by-scope</i> — proves consistent self-reference over an <i>exogenous</i> payoff, never reward-learning groundlessness. Exactly the primitive-reflective-oracle regime the no-go’s (R3)+(R5) premises exclude; listed for the same reason #deriv-self-actuation-grounding lists it: to mark the scope honestly.

Search Log:

- 2026-05-22 (*targeted, derivation-driven*): Cohen, Hutter & Osborne 2022 read directly (primary-source verified) across Assumptions 1–6, §”Competing models of the goal,” §”The assistance game,” §”Potential approaches,” and §”Conclusion.” Bareinboim-Correa-Ibeling-Icard 2022 CHT statement cross-checked; the floor (Lemma 1) verified to apply as a direct instance.
- 2026-05-22 (*targeted, prior-art adjacency*): The sister relationship to #disc-implementation-impossibility (Cohen 2022’s actor is the *agent*; GS/MS/Arrow’s actors are *designers*) is the formal placement carried by the five-row comparison in Discussion. Skalse, Howe, Krashennnikov & Krueger 2022 flagged as a separate #disc-implementation-impossibility charter-instance candidate, not coupled here.
- 2026-05-22 (*intuition-only on extensions*): Whether the unification with Result G’ via the single-interface commitment of V_{O_t} has been articulated elsewhere as a framework-level convergence is open. A targeted search on whether “ V_{O_t} information-narrowness as the engine behind both within-model self-modification no-gos and

across-model reward-learning no-gos” appears in the AI-safety / decision-theory literature is recommended before this segment advances past draft.

Working Notes.

- **Provenance.** Authored 2026-05-22 via Track CR Phase 4b per the integration-reconciliation verdict (spikes/spike-integration-reconciliation-2026-05-22/99-verdict.md §5 Phase 4b) and the Cohen-2022 strengthening verdict (spikes/spike-cohen-2022-strengthening-2026-05-22/99-verdict.md §3.b Option (ii); §4 cross-segment-ripple plan). Slug deriv-reward-channel-learning-no-go decided 2026-05-22 by Joseph (over the symmetric alternative deriv-goal-learning-grounding; *“I’m ok breaking the mold a bit to be more clear”*). The strengthening spike’s 02-formalization.md carries the full step-by-step derivation; this segment carries the canonical exposition.
- **Open before candidate.** (a) Targeted prior-art search on “ V_{O_i} information-narrowness as the engine behind two complementary no-gos” surfaced in the search history above; (b) the assistance-game extension (Cohen 2022 §“The Assistance Game”) formalization in AAT (currently scoped out by (R3)); (c) widening the EU-max scope to AAT’s broader strategy machinery — quantization, satisficing, lexicographic — by re-deriving the corollary; (d) the Class-3-coupled-substrate extension (currently flagged as open in Epistemic Status); (e) primary-source verification of Bareinboim-Correa-Ibeling-Icard 2022 (the CHT external anchor here is from secondary citation in #disc-identifiability-floor Instance 1 + the strengthening spike’s reading).
- **Cluster’s peer meta-segment home landed (2026-05-22, Track B cycle commit 5).** The deferred meta-promotion question from the Cohen-2022 strengthening verdict’s §2 step 4 is closed: the agent-side value-functional-grounding cluster has its own peer meta-segment home at #disc-value-functional-grounding-floor — sister to #disc-identifiability-floor (agent-side data-inference) and #disc-implementation-impossibility (designer-side mechanism-design). This segment is *Instance G* of that meta-segment (across-model failure; CHT-at-reward-channel engine; principal-side substrate-level commitment terminal grounding route via (R1)-(R5) escape menu). The two-routes-exhausts geometry (agent-side adaptive-substrate per Result G’ Corollary 2 + principal-side protocol-commitment per the escape menu here) is the meta-segment’s headline content. The #disc-constructive-impossibility-posture catalog now lists F (Result G’) and G (this segment) under the cluster’s home; the catalog landed pre-meta-segment per Track CR Phase 4b commit fcfb7a1, and the intro paragraph there updated per Track B commit 6 to name all three peer meta-segments cleanly.
- **Skalse 2022 as cross-cluster bridge candidate.** Skalse-Howe-Krashennikov-Krueger 2022 is structurally a designer-side no-go (no non-trivial unhackable reward-proxy pair on open policy sets) — naturally a #disc-implementation-impossibility candidate; the agent-side dual (agent exploiting the proxy) is naturally a member of this segment’s cluster. The Track B charter as scoped (GS / MS / Arrow) stayed at three per the strengthening verdict §2; the Skalse extension is therefore a *cross-cluster bridge evaluation*, not a single-cluster landing. Indexed at spikes/PROPOSED.md Tier 2 (added 2026-05-24) → spikes/PROPOSED-ADVANCED.md §Phase-3 #17 (two-pass evaluation: designer-side against #disc-implementation-impossibility’s filter; agent-side against this cluster’s). Decision: charter instance of one cluster, bridge segment between both, or honestly-adjacent-not-instance.

Incidental audit gold (lift 2026-05-31). Orthogonal pedagogical material from the de-novo auditors’ working dirs (per de-novo-audit-instructions.md §7.15). This segment (authored 2026-05-22) postdates most of the de-novo audit cycles, so it has no dedicated per-segment reflection; the single light note below was harvested from AUDIT-WORKING-773921’s reflection on the sister within-model no-go #deriv-self-actuation-grounding. No certified-track off-ramp.

Candidate Discussion / field contribution.

- **The two-no-go pairing as an alignment punchline.** A reader of the within-model no-go (#deriv-self-actuation-grounding) anticipated this across-model companion exactly: “if you can’t trust an agent to safely rewrite its own goals, [the sister segment] proves you can’t even trust an agent to correctly infer your goals from a reward channel.” The pair reads cleanly as: goal-stability fails *from the inside* (convention-monotonicity) and *from the outside* (CHT-at-the-reward-channel), both because V_{O_i} is too narrow to anchor it — so safety must be a substrate-level commitment, not a reward-channel inference. A candidate framing for the cluster’s shared headline (Claude, AUDIT-WORKING-773921).

Appendix F

Discussion: The Sandbox Evaluation Ceiling

A structural negative result about pre-deployment evaluation follows once the loop-as-Level-2-engine result (7.3) is in hand: a sandbox’s forkability makes it a *good* interventional laboratory — by overriding the agent’s action one obtains genuine Level-2 evidence *about the sandbox* — but evidence about the sandbox transfers to deployment only when the sandbox and deployment run on the same mechanisms. The ceiling on using sandbox evidence for deployment claims is therefore **transportability** / **external validity**: when the deployment causal system differs from the sandbox causal system at mechanisms on the relevant paths, the deployment intervention-response is *not transportable* from sandbox data alone, regardless of how thorough the evaluation is. This is an instance of the framework’s constructive-impossibility posture (5.6): the no-go is paired with the precise statement of what *would* be required to escape it — an invariance argument across the sandbox/deployment selection diagram, or deployment-time data.

Formal Expression.

A sandbox induces a structural causal model $\mathcal{M}_{\text{sb}}$ over a source domain π_{sb} ; deployment induces $\mathcal{M}_{\text{dp}}$ over a target domain π^* . The two share a causal diagram G , and their **selection diagram** D (Bareinboim & Pearl 2014, Def. 1) augments G with a square selection node $S_i \rightarrow V_i$ at every mechanism that may differ between the environments — that is, wherever $f_i^{\text{sb}} \neq f_i^{\text{dp}}$ or $P^{\text{sb}}(U_i) \neq P^{\text{dp}}(U_i)$. The quantity a deployment safety claim is about is the target interventional effect $R = P^*(o \mid \text{do}(a))$.

Derived (conditional on the sandbox/deployment selection diagram).

1. **Sandbox forkability is what makes the sandbox a Level-2 laboratory.** Resetting, replaying, and branching let the evaluator override the agent’s action — force a and observe the response — which samples the *mutilated* submodel $\mathcal{M}_{\text{sb},a}$ and yields genuine interventional data $P_{\text{sb}}(o \mid \text{do}(a), \mathbf{z})$ in π_{sb} (Bareinboim, Correa, Ibeling & Icard 2022, Def. 5; randomized controlled trials and A/B experiments are the canonical forkable realization of $\text{do}(\cdot)$, CHT footnote 19). Forkability is the source of Level-2 sandbox evidence, not a demotion of it.
2. **The target effect is identifiable from sandbox data iff it is transportable across D .** R is recoverable from $\{P_{\text{sb}}(o \mid \text{do}(a), \mathbf{z})\}$ exactly when it is *mz*-transportable in the sense of Bareinboim & Pearl 2014 (Thm. 1): $P^*(o \mid \text{do}(a), \mathbf{S})$ reduces, by the rules of do-calculus, to an expression in which the selection nodes $\mathbf{S}$ appear only on factors that are invariant between source and target.
3. **No-go.** There exist selection diagrams for which R is *not* transportable (Bareinboim & Pearl 2014, Fig. 1b — the smallest diagram exhibiting non-transportability even with X randomized in the source; the explicit two-model witness, Eqs. 3–4): two deployment models consistent with *all* sandbox data assign R different values. For these,

$$R = P^*(o \mid do(a)) \text{ is not identifiable from } \{P_{sb}(o \mid do(a), z)\}, \quad (F.1)$$

regardless of how thorough the sandbox evaluation is. This is the ceiling: when the deployment differs from the sandbox at mechanisms on the relevant paths — selection nodes on the wrong edges — no transport formula exists, and source-domain interventional evidence, however exhaustive, leaves the deployment effect undetermined.

1. **Boundary (positive content).** The no-go comes with an explicit characterization of what survives. Effects routed only through mechanisms invariant across D (no relevant S -node) *transport* — and these include both Level-1 and Level-2 sandbox claims; a perfectly faithful sandbox (no relevant selection nodes — deployment is mechanism-identical) transports completely. Effects routed through differing mechanisms (selection nodes on relevant paths) do *not* transport, and require either an invariance argument (an explicit claim that the relevant mechanisms agree across D) or deployment-time interventional data. The boundary is the selection diagram itself.

Epistemic Status.

Discussion-grade. The no-go is a direct application of two external results — the Causal Hierarchy Theorem (Bareinboim, Correa, Ibeling & Icard 2022) for the data-character side and transportability theory (Bareinboim & Pearl 2014) for the cross-environment side — to the AI sandbox/deployment setting; it introduces no new mathematics beyond constructing the selection diagram for that setting. It is *exact* conditional on the sandbox/deployment selection diagram: the no-go holds whenever the deployment SCM may differ from the sandbox SCM at mechanisms on the relevant paths, and dissolves when the selection diagram carries no relevant selection nodes (a mechanism-identical deployment, where the effect transports completely). The claim is conservative in form: it is about *identifiability of the deployment interventional distribution from sandbox interventional data*, not about any particular evaluation method’s measured accuracy. What is AAT-distinctive is not the transportability theorem but its placement — locating the much-debated “alignment evaluations don’t predict deployment behavior” pattern as an *external-validity / transportability* problem with an explicit selection-diagram boundary, rather than a measurement-quality, coverage, or model-capability one.

Discussion.

The negative result has positive content. It sharpens *which* safety claims a sandbox can and cannot establish. A sandbox is a genuine interventional laboratory — its forkability lets an evaluator force actions and measure responses, yielding real Level-2 evidence *about the sandbox*. What that evidence establishes about deployment is governed by the selection diagram: effects flowing only through mechanisms shared and invariant between sandbox and deployment carry over; effects flowing through mechanisms that differ do not, and for those, no amount of sandbox testing identifies the deployment effect. The framework’s prescription is therefore not “sandboxes are useless” but “sandboxes have a *structural* ceiling at the sandbox/deployment boundary, and deployment-time monitoring is not substitutable by pre-deployment evaluation when the relevant mechanisms differ across that boundary.” For a practitioner running production agent loops, this is the same content as: the monitoring done *after* deployment is doing structural work the pre-deployment evaluations cannot do, precisely on the mechanisms where the two environments diverge.

Why this is a ceiling, not a gradient. A measurement-quality limitation closes as rigor increases; a coverage gap closes as the test suite grows; a capability surprise closes as models are better understood. The sandbox ceiling closes under none of these, because the obstruction is external validity — whether the sandbox SCM and the deployment SCM agree on the relevant mechanisms — not the quantity or quality of sandbox data. More sandbox rigor produces more and better *sandbox* interventional evidence, which transportability theory still does not let one transfer to deployment across an adverse selection diagram. The only things that close the gap are an invariance argument across the selection diagram or deployment-time data — neither of which is “more sandbox testing.”

The downstream prescriptive content — what deployment-time monitoring must structurally look like to discharge what sandbox evaluation cannot — lands in the framework’s interventional-diagnostic / sidecar-monitoring treatment (the bi-predictability anchor at 89% vs 44% detection accuracy is developed in 7.2’s Discussion). The sandbox ceiling is what gives that downstream monitoring work its load-bearing role: it is not a redundant safety layer but the only layer with access to the Level-2 quantity the safety claim is actually about.

Findings.

The Sandbox Evaluation Ceiling. Brief: A sandbox lets you reset and replay a system — and that replayability is what makes it a *good* laboratory: you can force an action and watch the response, which is genuine interventional (Level-2) evidence *about the sandbox*. The catch is not the sandbox’s data type; it is that the sandbox world and the deployment world may run on different mechanisms. There is a theorem (Bareinboim–Pearl transportability) saying that whether a sandbox-measured effect carries over to deployment depends on *where* the two worlds differ: effects flowing only through shared, invariant mechanisms transport; effects flowing through mechanisms that differ do not — and for those, no amount of sandbox testing identifies the deployment effect. So the ceiling on pre-deployment evaluation is *external validity*, and the missing piece is either an explicit invariance argument or deployment-time data — which is why runtime monitoring is structurally non-substitutable, not redundant.

Impact: Gives the much-debated observation that “alignment evaluations don’t predict deployment behavior” a precise *structural* mechanism — external validity across a selection diagram — rather than a measurement-quality, coverage, or capability story, together with an exact statement of where the boundary lies (effects through mechanisms invariant across the sandbox/deployment selection diagram transport, including Level-2 effects; effects through differing mechanisms do not, regardless of evaluation thoroughness). The constructive flip is the operational payload: deployment-time monitoring is not a redundant safety layer but the only channel with access to the deployment interventional effect when the relevant mechanisms differ across the boundary, which is why pre-deployment evaluation cannot substitute for it — and the repair is named precisely (an invariance argument across the selection diagram, or deployment-time data).

Novelty Claim: *Application of established machinery* (Bareinboim & Pearl 2014 transportability / selection diagrams, resting on the Pearl 2009 / Bareinboim, Correa, Ibeling & Icard 2022 do-calculus and Causal Hierarchy Theorem) to the AI sandbox/deployment evaluation gap. The transportability apparatus is adopted, not invented; the AAT-distinctive contribution is locating “pre-deployment AI evaluations don’t predict deployment behavior” as a *selection-diagram transportability* problem — with the explicit transportable-vs-not boundary as a function of where the sandbox and deployment SCMs differ, and the constructive flip that runtime monitoring is the only channel with access to the deployment effect on the differing mechanisms. The framing reclassifies a whole family of pre-deployment safety-evaluation approaches as facing a structural external-validity ceiling rather than a closable coverage gap.

Related Work:

Table F.1

<i>ASF concern</i>	<i>Prior-art language</i>	<i>Relationship / positioning</i>
Cross-environment generalization of interventional effects	Bareinboim & Pearl 2014, <i>Transportability from Multiple Environments with Limited Experiments</i> (selection diagrams, <i>mZ</i> -transportability)	<i>formal antecedent</i> — supplies the selection-diagram apparatus and the transportability no-go (a source-domain interventional effect need not transport to a target domain); AAT supplies the construction of the sandbox/deployment selection diagram

<i>ASF concern</i>	<i>Prior-art language</i>	<i>Relationship / positioning</i>
Interventional vs observational identifiability	Pearl 2009 <i>Causality</i> 2nd ed.; Bareinboim, Correa, Ibeling & Icard 2022 (Causal Hierarchy Theorem)	<i>machinery</i> — grounds the data-character side (a forkable harness that overrides the action yields genuine Level-2 sandbox data) and the do-calculus the transportability reduction runs on
Evaluation–deployment gap in AI safety	Practitioner literature on eval/deployment mismatch (qualitative)	<i>conceptual precursor</i> — names the phenomenon without a structural mechanism; AAT supplies the transportability mechanism and the exact selection-diagram boundary

Search Log: - 2026-05-30 (*intuition-only* on the application-novelty): transportability theory and selection diagrams (Bareinboim & Pearl 2014) are established and adopted directly. The unsearched claim is whether the specific framing — *the AI pre-deployment-evaluation / deployment gap is a selection-diagram transportability problem, with an explicit transportable-vs-not boundary as a function of where the sandbox and deployment SCMs differ* — has been stated elsewhere as a structural ceiling on pre-deployment evaluation. Pre-search expectation: the constituent move (constructing a selection diagram for a source/target pair and reading off transportability) is standard within transportability theory; the framing as a *structural* ceiling on the AI-safety evaluation program, with the constructive flip to runtime monitoring as the only channel with deployment-effect access, is plausibly AAT-distinctive but not yet verified under nominally-comprehensive search. Targeted future candidates: AI-evaluation / model-evaluation literature on distribution shift vs. interventional generalization; causal-transportability work applied to sim-to-real or offline-to-online RL deployment guarantees.

Working Notes.

- Promoted to its own segment 2026-05-28 (PROPOSALS SP-25; audit-773921 Finding 4, Gemini). The content previously lived only as a ## subsection inside the chapter-end discussion 7.10 (still cross-referenced from there). The promotion was an author-and-auditor convergence: the segment author's own Working Notes in 7.10 had flagged it as a promotion candidate pending a future cycle, and the de-novo auditor independently surfaced it as a buried high-impact result.
- Open placement question carried from 7.10's Working Notes, now resolved toward Appendix A (the constructive-impossibility no-go home, alongside B and C) rather than an extension of 4.10's Discussion — keeps the no-go discoverable in the family it belongs to and lets 5.6 reference it as an instance.
- The theorem-grade formalization of the underlying loop-Level-2 result lives in NeurIPS 2026 Paper 2 (~/src/neurips/02-unified-convergence-rl/, §5 / the (C1)/(C2)/(C3) triple); that is an in-review external draft, not canon, so the pointer is kept here in Working Notes rather than in the segment body. When that work is published it becomes an external citable object the body can cite directly.
- **Re-grounding 2026-05-30 (transportability spike spikes/.integrated/spike-causal-access-transportability-2026-05-30.md).** The segment originally stated the ceiling as a Pearl-level demotion ("forkable execution yields Level-1 data; deployment is Level-2; the Causal Hierarchy Theorem forbids the crossing"). A de-novo auditor (AUDIT-WORKING-526815) challenged the forkability→Level-1 step, and the strengthen-first spike verified it false against the primary sources: forkability is orthogonal to the L1/L2 distinction (CHT Def. 5 — the *do*-semantics is population-level; CHT footnote 19 — RCTs/A/B tests are forkable *and* the canonical source of Level-2 data), so the discriminator is mechanism-override, not replayability. The strengthening (not a retraction): the correct ceiling is *transportability / external validity* (Bareinboim & Pearl 2014 selection diagrams), which is sharper — it preserves the no-go (Fig. 1b non-transportable witness), names which claims transport (mechanism-invariant ones, Level-1 *and* Level-2 alike) and which do not, and names the repair (an invariance argument across the selection diagram, or deployment-time data). The former candidate-strengthening note (state forkability-implies-Level-1 as its own lemma) is deleted: that lemma is false; the transportability demonstration in the Formal Expression is the replacement strengthening.

Incidental audit gold (lift 2026-05-31). Swept in the gold-lift pass (per de-novo-audit-instructions.md §7.15): **no incidental pedagogical/framing gold found.** This segment was promoted to its own file only 2026-05-28, so it postdates the de-novo audit cycles and has no dedicated per-segment reflection. The one auditor thread that touched its content (AUDIT-WORKING-526815's forkability-implies-Level-1 challenge) was a *certified-track* finding, already strengthen-first-resolved into the transportability re-grounding recorded in the note above — not gold. Marker left so future sweeps know this segment was checked.

Appendix G

Result: Sector-Persistence Template

The abstract Lyapunov persistence argument *stated once, parameter-free*, so that the framework’s many persistence-flavored results can be cited as *instances* without re-proving the argument. Any state variable ξ evolving under bounded-correction dynamics with bounded disturbance admits the same persistence conclusion under three preconditions: **(T1)** zero correction at zero state (no correction when the state is at target); **(T2)** local sector condition (the correction function points inward with at least baseline efficiency α in a radius- R region); **(T3)** bounded disturbance (deterministic bound, or stochastic with bounded variance). Under these preconditions the template gives the chapter’s central inequality form: the state is ultimately bounded by *disturbance rate / sector parameter* (Model D, $1/\alpha$ scaling) or *root-of-variance times root-of-dimension over sector parameter* (Model S, $1/\sqrt{\alpha}$ scaling). The two scalings produce the framework’s Model D vs Model S distinction that propagates through every result that cites the template.

The template’s value is *taxonomic economy*. AAT’s persistence-flavored results — epistemic, strategic, team, composite closure, composite tempo, adversarial destabilization, and identity-continuity across turnover — are instances of one template. Each instantiation specifies only what varies: its state variable, correction function, effective disturbance rate, and reserve. The Lyapunov boilerplate is absorbed once; each downstream result’s distinctive content is just the *effective-disturbance decomposition* particular to its setting. Persistence and destabilization are the same inequality viewed in opposite directions — B ’s persistence certificate holds iff $\alpha_B R_B > \rho_B$, and is lost (with escape forced in the radially tight case) iff $\rho_B \geq \alpha_B R_B$ — and the superlinear adversarial scalings ($b = 2$ Model D, $b = 3/2$ Model S) follow without further derivation. The mathematical home is monotone-operator theory (Rockafellar 1970; Bauschke-Combettes 2017): **(T2)** is a *one-point strong monotonicity* condition at the equilibrium $\xi^* = 0$ — strictly weaker than the full two-point strong monotonicity the bridge lemma requires for composition closure. AAT’s relationship to that literature is *specialization*, not generalization; acknowledging this is what makes the template feel *earned* rather than parochial.

Formal Expression.

Let $\xi(t) \in \mathbb{R}^n$ be a state variable evolving under

Template preconditions (sector-persistence-template).

$$\frac{d\xi}{dt} = -F(\xi) + w(t) \tag{G.1}$$

where F is a correction function and $w(t)$ is a disturbance. The template applies when:

(T1) Zero correction at zero state. $F(0) = 0$ — no correction is applied when the state is at its target.

(T2) Local sector condition. There exist $\alpha > 0$ and $R > 0$ such that

$$\xi^T F(\xi) \geq \alpha \|\xi\|^2 \quad \text{for } \|\xi\| \leq R. \quad (\text{G.2})$$

The correction points inward with at least baseline efficiency α throughout the region of radius R .

(T3) Bounded disturbance. Either:

- *Model D (deterministic bound):* $\|w(t)\| \leq \rho_\xi$, or
- *Model S (stochastic zero-mean):* $\mathbb{E}[\|w(t)\|^2] = \sigma_\xi^2$ with $w(t)$ a Wiener-process increment.

Under (T1)–(T3), with $V(\xi) = \frac{1}{2}\|\xi\|^2$ as Lyapunov function:

Template result (from sector-condition-derivation Props A.1, A.1S, A.2).

Model D. The state is ultimately bounded by $R^* = \rho_\xi/\alpha$. Structural persistence (the ultimate bound fits within the sector-condition region) requires

$$\alpha > \frac{\rho_\xi}{R}. \quad (\text{G.3})$$

The adaptive reserve — the additional disturbance the system can absorb before persistence fails — is $\Delta\rho_\xi^* = \alpha R - \rho_\xi$.

Model S. The state satisfies $\mathbb{E}[\|\xi(t)\|^2] \rightarrow n\sigma_\xi^2/(2\alpha)$ in mean square, giving RMS bound $R_S^* = \sigma_\xi\sqrt{n/(2\alpha)}$. Structural persistence in the mean-square sense requires

$$\alpha > \frac{n\sigma_\xi^2}{2R^2}. \quad (\text{G.4})$$

The Model D result scales as $1/\alpha$; the Model S result scales as $1/\sqrt{\alpha}$ — correction is less effective against noise than against drift. This scaling difference propagates into the adversarial exponent regimes (16.3): $b = 2$ under Model D, $b = 3/2$ under Model S.

Instantiations in AAT. The template is invoked across seven segments. Each specifies its own (ξ, F, ρ_ξ, R) and verifies (T1)–(T3) locally:

Table G.1

<i>Segment</i>	ξ	<i>Effective ρ_ξ</i>	R	<i>Locally-verified precondition</i>
4.5	δ_t (epistemic mismatch)	ρ (environmental disturbance rate)	model-class capacity, or task-adequacy threshold $\ \delta_{\text{critical}}\ $ if stricter	(T2) via 4.3
9.10	δ_Σ (strategic mismatch)	ρ_Σ (rate of edge invalidation)	R_Σ (strategic reserve)	(T2) via Beta-Bernoulli edge updates (R Props B.1–B.6); constant- α requires experience discounting

<i>Segment</i>	ξ	<i>Effective</i> ρ_ξ	R	<i>Locally-verified precondition</i>
15.1	δ_l (sub-agent mismatch)	$\rho_l^{\text{eff}} = \rho_{l,\text{env}} + \sum_j \gamma_{j \rightarrow l}^{\text{adv}} \mathcal{T}_j - \sum_j \gamma_{j \rightarrow l}^{\text{coop}} \mathcal{T}_j$	R_l	Cooperative coupling can drive ρ_l^{eff} below the single-agent ρ
13.1 (bridge lemma)	e_m (trajectory error at macro-boundaries m)	$\varepsilon^* \nu_c$ (closure-defect per macro-step $\times$ macro-update rate)	the composite's R_c	Tier-specific contraction stronger than (T2): incremental sector bound (DA2'-inc). Tier 1 proved; Tier 2 local; Tier 3 domain-specific
13.2	δ_c (composite mismatch)	$\rho_{\text{ext}} + \varepsilon^* \nu_c$ (external + internal)	R_c	(T3) requires the bridge-lemma contraction; $C_{\text{coord}} \geq \varepsilon^* \nu_c$ is the tempo-equivalent of the internal disturbance
15.2	δ_B (target-agent mismatch)	$\rho_{B,\text{base}} + \gamma_A \mathcal{T}_A$ (Model D) or $\sigma_{B,\text{base}} + \gamma_A \mathcal{T}_A$ (Model S)	R_B	Destabilization is the negation of persistence: (T3)'s coupling-amplified disturbance violates $\alpha R > \rho_\xi$
#der-identity-continuity-threshold	g_k (identity gap, turnover-indexed)	ρ_k (projected static identity rate-distortion floor, per boundary)	D_Δ (fidelity-bounded identity-relevant information)	The reflected (Lindley) discrete instantiation on the turnover axis; compensation is relational re-grounding $\mathcal{G}_{\text{rg},k}$; Model-S-family, $\mu = 0$ driftless boundary load-bearing; conditional under (M-ADD)/(M-FREE)/(C-S) — the $\mathcal{A}_{\text{refl}}$ operator family (see Discussion)

The same Lyapunov argument applies in every row. What varies is how each instantiation defines its state variable, its correction function, and — most importantly — what counts as “effective disturbance” for its context (environmental noise, adversarial coupling, closure defect, or a decomposition of these).

External mathematical lineage: monotone-operator theory. The sector-Lyapunov apparatus this segment factors out has a well-established mathematical home: AAT’s sector condition (T2) is a **one-point strong monotonicity** condition in the sense of Rockafellar 1970 (*Convex Analysis*, §24) anchored at the equilibrium $\xi^* = 0$, and the incremental strengthening DA2'-inc required by #form-composition-closure’s bridge lemma is **full two-point strong monotonicity** in the Bauschke-Combettes 2017 (*Convex Analysis and Monotone Operator Theory in Hilbert Spaces*, 2nd ed., §§22–28) monotone-operator framework. The Lyapunov-plus-Grönwall argument proving the Model D ultimate bound and the Itô-Lyapunov argument proving Prop A.1S are standard specializations of the **monotone-operator perturbation theorem** (Bauschke-Combettes Thm 5.14–5.16; Parikh & Boyd 2014, *Proximal Algorithms*, §2). Banach-Picard contraction under bounded perturbation, monotone-operator convergence under square-summable noise, operator-splitting for composite systems (Douglas-Rachford, ADMM) — the full supporting apparatus is available in that literature.

AAT’s relationship to monotone-operator theory is *specialization + repurposing*, not generalization. AAT-distinctive content: (i) **one-point anchoring at the equilibrium** — strictly weaker than full two-point strong monotonicity, matched to fixed-point-at-target semantics, admitting agent classes (PID-bounded-plant, variational-approximate) where

full monotonicity fails but persistence-at-the-target is available. (ii) The **Model D / Model S disturbance decomposition** with distinct $1/\alpha$ vs $1/\sqrt{\alpha}$ scaling — monotone-operator theory has perturbation theorems but no systematic bounded-adversarial vs stochastic-zero-mean split propagating to adversarial-exponent regimes at $b = 2$ vs $b = 3/2$. (iii) Composition with the `#disc-identifiability-floor` meta-pattern supplies a second (information-theoretic) axis orthogonal to the operator-theoretic machinery. (iv) The `#disc-composition-consistency` cross-level commitment operates at AAT’s level of description (agent / composite / macro-agent) rather than at the abstract operator level. (v) The sub-scope α/β epistemic labeling is scope-honesty, not a mathematical partition — it tracks which agent classes give the monotone-operator structure *by construction* versus *by per-instance verification*.

This acknowledgment is load-bearing for scope honesty: the mathematical machinery is established external to AAT; AAT’s distinctive content is the agent-architecture specialization (singular-trajectory identity per `#scope-agent-identity`; signed-coupling composition; coordinate-forcing via uniqueness theorems per `#disc-additive-coordinate-forcing`; three meta-patterns) rather than novel monotone-operator mathematics. `#deriv-sector-condition`’s Grounding paragraphs name the specific operator-family correspondence (proximal / firmly-nonexpansive / cocoercive / strongly-monotone-gradient / linear-PD) for the five sub-scope- α agent classes. `#result-contraction-template` extends to non-Euclidean metrics via Lohmiller-Slotine differential-contraction (also within the broader monotone-operator lineage). The honest limits of the unification: the coarse-graining projection Λ (13.1) does not fit the operator-sector primitive (heterogeneous spaces, three independent admissibility conditions); three of five metric- α_2 cases in `#result-contraction-template` remain theorem-imported rather than AAT-internally derived; the identifiability-floor axis is orthogonal to the operator-sector axis.

Comparison with the FEP-flow stability argument. Active inference’s stability arguments come from the geometry of the variational free-energy landscape — agents are argued to flow toward the minimum of variational free energy on a non-equilibrium-steady-state (NESS) density. The primary source for the NESS-density framing is Friston 2019, “A free energy principle for a particular physics,” arXiv:1906.10184; the path-integral / particular-kinds methodological extension is Friston, Da Costa, Sakthivadivel, Heins, Pavliotis, Ramstead & Parr 2023, “Path integrals, particular kinds, and strange things,” *Phys. Life Rev.* 47 (which rewrites the FEP-flow argument in path-integral language rather than proving new stability bounds). Aguilera, Millidge, Tschantz & Buckley (2022, “How particular is the physics of the free energy principle?”, *Phys. Life Rev.* 40:24–50) showed that the FEP-flow argument’s mathematical validity is narrow: the NESS-density framing holds only in a small parameter regime for non-equilibrium linear stochastic systems, and natural extensions (nonlinear, non-Gaussian, non-equilibrium) often fall outside the proven regime.

The AAT persistence template is structurally different: it is a Lyapunov-based argument requiring only (T1) zero-correction-at-zero-state, (T2) local sector condition (correction points inward), and (T3) bounded disturbance — all of which are checked locally for each instantiation (A Props A.1, A.1S, A.2). The template applies to bounded and to mean-square-stochastic disturbance, gives explicit ultimate-bound and adaptive-reserve formulas, and does not depend on NESS structure or on a free-energy gradient.

The breadth difference is not rhetorical: where the FEP-flow argument’s parameter regime is debated in the AI literature itself, the sector-Lyapunov apparatus is the standard machinery of nonlinear control theory (Khalil 2002, *Nonlinear Systems*, 3rd ed., Prentice Hall, ch. 4) and applies wherever (T1)–(T3) hold. This is one of AAT’s stronger structural positions and is worth making explicit when comparing AAT to active inference: AAT does the persistence work AI tries to do, with broader validity and explicit ultimate-bound formulas.

Epistemic Status.

Exact. The template is the abstract form of the Lyapunov result proved in A (Props A.1, A.1S, A.2). The proofs transfer without modification whenever (T1)–(T3) hold; the template’s contribution is the recognition that AAT’s persistence-flavored results are instances of a single pattern and the enumeration of what each instantiation must verify.

On (T2) and A2' sub-scoping. (T2) is the local sector condition A2' transcribed to a generic state variable ξ . Within the state-variable spaces relevant to AAT's instantiations (epistemic mismatch δ , strategic mismatch δ_Σ , sub-agent mismatch δ_i , composite trajectory error e_m , composite mismatch δ_c , target-agent mismatch δ_B), the update rules are overwhelmingly sub-scope α in the 4.2 / 4.3 sense — Bayesian, exponential-family, strongly-convex-gradient, or linear-PD. For these, (T2) is *derived* from the update-rule structure under B1 directional fidelity. **Important, and easy to miss: this is not what carries forward universally.** For sub-scope β instantiations (e.g. a team with a rule-based sub-agent) (T2) is an *empirical precondition, verified per-instantiation*, not derived. That distinction scopes which instantiations satisfy (T2) by which route; it does not touch the template's exactness, which is the conditional “(T1)–(T3) $\Rightarrow$ persistence” — anything satisfying (T2) by either route inherits the exact result.

On Prop A.1S region-awareness. The Model S case uses the region-aware form of Prop A.1S (stopped second-moment bound + the fixed-time tail $P(\|\xi(t)\| > R) \leq \text{const}$ under the mean-square persistence condition). Instantiations that rely on the Model S result automatically inherit the stopping-time localization — no extra work is required at the template level. As at A, Model S provides a distributional / fixed-time guarantee, not pathwise-forever containment (which is a Model-D-only guarantee); template instantiations inherit that kind-of-guarantee distinction, not an infinite-horizon non-exit bound.

Max attainable: *exact*. The result is as strong as Lyapunov stability theory. Additional work could extend the template (state-dependent noise, time-varying α , non-quadratic Lyapunov functions) but would not strengthen it within its stated scope.

What the template does not establish:

- *Quantitative convergence rates.* Lyapunov gives stability and ultimate bounds, not convergence speed. Specific rates require instantiation-level analysis.
- *Behavior outside the sector-condition region.* (T2) holds only for $\|\xi\| \leq R$. Beyond R , correction may break down — the structural-adaptation regime of 4.8.
- *Time-varying parameters.* The template assumes constant α and R . Time-varying cases require additional machinery. 9.10 documents the most important AAT example: $\alpha_\Sigma = 1/(n + 1)$ decays with experience, so strategic persistence requires experience discounting at rate $(1 - \lambda) > \rho_\Sigma/R_\Sigma$ to recover constant- α .
- *Heavy-tailed disturbances.* Neither (T3-D) nor (T3-S) covers disturbances with unbounded moments. Extreme tail events are better handled as triggers for structural adaptation than as disturbances the template can absorb.

Discussion.

Why factor this out. Every persistence-flavored result in AAT takes the form “correction rate exceeds effective disturbance rate (relative to reserve)” and is proved by the same Lyapunov function $V(\xi) = \frac{1}{2}\|\xi\|^2$. Stating the argument once, parameter-free, makes visible that the theory is *one result about bounded-correction dynamics applied to several state spaces*, not six separately-proved results that happen to look alike.

Each instantiation's distinctive content is now sharply visible: *what is the state variable, and what counts as its effective disturbance?* The Lyapunov machinery is shared; the characterization of ρ_ξ is where the domain-specific insight lives. Adversarial destabilization's content is the coupling term $\gamma_A \mathcal{T}_A$; team persistence's content is the cooperative-minus-adversarial decomposition; composition closure's content is the closure-defect rate $\varepsilon^* \nu_c$. In each case, the template absorbs the Lyapunov boilerplate and lets the distinctive claim stand without it.

This template is the *interior facet* of the stability certificate (5): the “correction rate exceeds effective disturbance rate” condition is the certificate being positive-definite on the scope ball, and the existence of such a certificate is equivalent to exponential stability (H). The shared Lyapunov function $V(\xi) = \frac{1}{2}\|\xi\|^2$ is the certificate in the Euclidean metric; AI carries the non-Euclidean-metric interior.

The template is the typed-bridge shared upstream — two-model, not one operator. This template is the canonical anchor of the *accumulation typing* convention (NOTATION.md §"Accumulation typing"): the bridge from a per-step residue to its accumulated consequence *is* this conditional. The bridge is **two-model, not a single bounded operator** — (T3)'s Model D and Model S are co-equal, and the $1/\alpha$ vs $1/\sqrt{\alpha}$ scaling (line above) plus the Cor-A.1S.1 categorical containment dichotomy (A) are the structural fingerprint that they are *different functionals*, not one functional read in two norms (a change of α -homogeneity *degree* cannot come from a norm choice on a fixed operator). The across-turnover axis is the worked demonstration that the split is structural, with two operator families at *opposite ends of the same singular contraction parameter*, **neither superseding the other**: the $\mathcal{A}_{\text{ref}}$ family — a reflected (Lindley/Loynes) walk whose load-bearing content is its driftless $\mu = 0$ boundary — *instantiates* the template on the turnover index (#der-identity-continuity-threshold, the new row above); the $\mathcal{A}_D$ family — the linear destroy-and-reconstruct contraction whose affine norm *diverges* exactly as the contraction gap closes — is the **honest scope boundary where the template provably does not transfer** (#der-turnover-information-recursion: an unforced multiplicative information decay has no correction function pointing a perturbed state inward, so (T1)–(T3) have no counterpart; the non-transfer is a definite result, not a caveat). The $\mu = 0$ boundary of the $\mathcal{A}_{\text{ref}}$ instantiation *is* the singularity of the $\mathcal{A}_D$ regime — you cannot linearize across that pole, which is why the two are distinct operators and the accumulation-type confound (treating one as a normalization or linearization of the other) is a category error. The typing convention exists precisely so this — *which model, which operator family* — is carried in the notation rather than silently reconstructed.

What each instantiation must still verify. Invoking the template is not trivial. Each instantiation must establish (T1)–(T3) for its specific (ξ, F, ρ_ξ, R) , and the non-trivial verifications differ substantively:

- *Strategic persistence* (9.10): (T2) is verified for Beta-Bernoulli edge updates across five DAG topologies (R Props B.1–B.6). But $\alpha_\Sigma = 1/(n + 1)$ is time-varying: it decays monotonically with experience. Constant- α — and therefore the template's trajectory guarantee — requires experience discounting as a prerequisite, not a heuristic.
- *Composition closure* (13.1): (T2) as stated is insufficient. The bridge lemma requires the *incremental sector bound* (DA2'-inc) — strongly monotone F across the whole state space, strictly stronger than the one-point sector bound (T2). Three agent tiers result: Tier 1 where contraction is proved for the full class (Bayesian on exponential families, gradient descent on strongly convex losses, linear correctors with positive-definite gain), Tier 2 where it holds locally, Tier 3 where it must be verified per-domain.
- *Team persistence* (15.1): (T3) is the decomposition $\rho_i^{\text{eff}} = \rho_{i,\text{env}} + \sum_j \gamma_j^{\text{adv}} \mathcal{T}_j - \sum_j \gamma_j^{\text{coop}} \mathcal{T}_j$. Cooperative coupling enters through a negative term, which can drive ρ_i^{eff} below the single-agent ρ — this is formally how teams persist where individuals cannot.

Persistence and destabilization as one result. Adversarial destabilization (15.2) is the template applied with coupling-amplified disturbance: $\rho_B = \rho_{B,\text{base}} + \gamma_A \mathcal{T}_A$. The "destabilization threshold" — the condition under which agent A pushes agent B past its stability boundary — is precisely the *negation* of the template's persistence condition for B 's instantiation. B 's persistence certificate holds iff $\alpha_B R_B > \rho_B$; the certificate is lost — destabilization, with escape forced when B 's sector bound is radially tight — iff $\rho_B \geq \alpha_B R_B$. These are the same inequality viewed in opposite directions, not independent results; both directions are certificate-level statements, and the agent-level only-if requires tightness (A Lemma A.1N). The superlinear adversarial scaling ($b = 2$ under Model D, $b = 3/2$ under Model S) follows from the template's $1/\alpha$ vs $1/\sqrt{\alpha}$ scaling without further derivation.

Signed-coupling pattern across instantiations. The six instantiations above — plus J, which derives a closed-form composite (T2) for the matched-symmetric-Tier-1 dyad — share a signed-coupling structure: each effective ρ_ξ is a sum of environmental disturbance plus a cross-agent contribution whose sign encodes cooperative vs adversarial coupling. Team persistence, composition closure (via bridge-lemma), tempo composition, adversarial destabilization, and composite critical-mass are instances of one inequality at different state-variable levels. The template is where this pattern lives at the meta level; the per-segment instantiations supply the domain-specific (ξ, F, ρ_ξ, R) quadruples.

Cost complement: the template's information-rate floor. Alongside the threshold condition, each instantiation inherits an information-rate lower bound from I: under Model S with Gaussian-OU-shaped disturbance statistics on ξ , the sustained information rate required to maintain $\mathbb{E}[\|\xi\|^2]_{ss}$ at the ultimate bound satisfies $\dot{R} \geq n\alpha/2$ nats/time with the instantiation's own $(\alpha, n, \sigma_\xi^2)$. For the epistemic-mismatch instance, this translates to the channel-capacity floor $C \geq \mathcal{T}/2$ — a first-class persistence prerequisite that the threshold condition alone does not surface. Each instantiation can derive its own cost bound by direct substitution into the main theorem of I; the consolidation of this into a parametric template-cost subsection is flagged in that segment's Working Notes.

The coordination overhead C_{coord} is effective disturbance at the composite level. The tempo-composition inequality $\mathcal{T}_c \leq \sum_i \mathcal{T}_i$ has a lower bound on the gap: $C_{\text{coord}} \geq \varepsilon^* \nu_c$ (13.2). This is the template's instantiation with $\xi = \delta_c$ and effective disturbance $\rho_{\text{ext}} + \varepsilon^* \nu_c$ — the composite's internal closure error acts as an additional disturbance at the macro level, absorbing correction capacity that would otherwise address the external environment. Brooks's Law follows as an instance: adding agents increases $\sum_i \mathcal{T}_i$ but may increase $\varepsilon^* \nu_c$ faster, pushing ρ_c^{eff} above $\alpha_c R_c$.

Relationship to 4.5. 4.5 is the canonical single-agent instantiation. It carries additional content beyond the template — *task adequacy*, a domain-specific constraint $R^* < \|\delta_{\text{critical}}\|$ that further restricts the usable region — and it instantiates the template in the linear special case where $F = \mathcal{T}\delta$, $\alpha = \mathcal{T}$, and (T2) holds globally. Task adequacy is not part of the template because it is not part of the Lyapunov argument; it is an application-level constraint the agent's domain imposes.

Relationship to 11. The composition-consistency postulate requires that AAT's predictions be compatible across levels of description. The template is how this compatibility cashes out operationally: the same persistence argument applies at every level where a state variable with a sector-bounded correction function is present. Part III segments that invoke the template (team persistence, composition closure, tempo composition) are applying AAT's single persistence argument at the composite level, with effective disturbance decompositions that capture what is distinctive about the composite scope.

Relationship to AI (metric-formulation generalization). The Euclidean sector inequality (T2) is the $M = I$ specialization of a broader contraction-metric condition (CT2) under a Riemannian metric M (Lohmiller & Slotine 1998). For several AAT-relevant agent classes, the natural Lyapunov lives in a non-Euclidean metric — Fisher for statistical-manifold learning (matrix Kalman in information metric; exponential family in natural parameters), Hessian-induced for ill-conditioned strongly-convex optimization, Lyapunov-equation-determined for asymmetric-stable linear systems, or Lyapunov-metric for PID-with-bounded-plant-nonlinearity. AI states the generalization once with (CT1)–(CT3) preconditions matching (T1)–(T3), adds compositional theorems (parallel / cascade / negative-feedback with small-gain) that extend #deriv-critical-mass-composition (CM2) to heterogeneous sub-agents via (CM2-M), and fills #disc-separability-pattern's seventh ladder (A2'-scope into metric- α_1 / metric- α_2 / metric- β). The Euclidean formulation stated in this segment remains the default for Euclidean-natural instances; AI is invoked when the natural coordinate is non-Euclidean. Structural consequence worth noting: the (CT2) condition at $M = I$ is equivalent to #form-composition-closure's DA2'-inc (incremental sector bound), so AAT has been carrying the Jacobian-level Euclidean contraction condition at the composite level all along; AI makes this explicit at the single-agent level.

Working Notes.

- The template is type result because it states a result abstractly; the proof lives in A (the canonical instance for δ), and the transfer to other state variables is routine once (T1)–(T3) are verified. Treating it as a formulation would misrepresent its status — the Lyapunov argument does not leave room for alternative formulations within its stated scope.
- Candidate extensions that would strengthen the template: (i) time-varying $\alpha(t)$ with lower bound, for the strategic case; (ii) state-dependent noise intensity, for cases where the disturbance magnitude depends on the current mismatch; (iii) non-quadratic Lyapunov functions, for richer stability regions. Each is a real extension of Lyapunov theory, not specific to AAT.

- The template’s role in 11’s “applies at every level” claim is now explicit: the claim holds at every level where (T1)–(T3) can be verified for the level’s state variable. This sharpens the composition-consistency commitment from “applies at every level” to “applies at every level the template applies to,” which is a testable condition rather than a universal assertion.
- **Dynamic-regime placement (added 2026-05-21 per Phase 5 cross-segment ripple).** The template’s (T1)–(T3) preconditions hold cleanly on **R0 contraction-regime** composites per #disc-dynamic-regime-axis; the sector constant α and basin radius R live at the single-attractor structure that R0 supplies. The template lifts to **R1 equilibrium-regime** composites via the A2’-analog transfer in #deriv-strategic-composition (sub-scope α'): the state variable becomes $\xi = \pi - \pi_*$ (deviation from Nash) and α_{joint} lives at the joint potential’s curvature (potential games, Monderer-Shapley 1996) or the joint Jacobian’s symmetric part (monotone games, Rosen 1965) — see #disc-dynamic-regime-axis §”Lyapunov-machinery transfer R0 $\rightarrow$ R1” for the explicit form. **R2 cyclic-distributional-regime** composites (sub-scope β') do *not* admit the sector-template — the state-space Lyapunov form fails; CCE convergence at $O(1/\sqrt{T})$ via Hart-Mas-Colell 2000 is the replacement machinery. Any cite of this template’s (T1)–(T3) preconditions should name the regime tier the cited claim sits at.

Incidental audit gold (lift 2026-05-31). Orthogonal pedagogical / generative material from the de-novo auditors’ working dirs (per de-novo-audit-instructions.md §7.15), deduplicated and lightly attributed. Substrates consistently called this segment the framework’s “central nervous system” / “architectural spine” — the meta-theorem the whole core hangs on. Certified-track item off-ramped at the end.

1. Candidate Briefprose / pre-prose.

- **The one-inequality unification.** “It unifies learning, teamwork, strategy, and adversarial deception under a single inequality: $\alpha R > \rho_\xi$. If your efficiency $\times$ capacity exceeds your effective disturbance, you persist.” A compact, accurate statement of what the template buys (Gemini, AUDIT-WORKING-849201 and AUDIT-WORKING-193847).
- **“Survival checklist” framing.** The template reads as a universal checklist: for *any* adaptive system, define its state variable, verify it corrects toward zero (T1), verify the correction is sector-bounded (T2), bound its disturbance (T3) — and the math then guarantees survival (Gemini, AUDIT-WORKING-193847).

2. Candidate Discussion.

- **“Combat is not special math — it is just $w(t)$.”** Adversarial destabilization is the *negation* of the persistence template: an attacker is simply another agent acting as the disturbance term in your survival equation. There is no separate “combat module” — surviving a harsh winter and surviving a targeted attack are mathematically identical operations, both requiring $\alpha R > \rho$. The corollary: “defense” is not separate from “learning”; the same correction function F that learns physics is the one that evades a predator. *Epistemic integrity is armor* (Gemini, AUDIT-WORKING-193847). A striking framing for the segment’s role as the common root of cooperative/adversarial dynamics.
- **The Active-Inference contrast, sharpened.** The “external mathematical lineage” subsection’s contrast with Friston’s active inference is read as devastatingly clean: AI relies on fragile non-equilibrium-steady-state density assumptions, whereas AAT relies on robust standard nonlinear control theory (Bauschke-Combettes monotone-operator grounding). Candidate to surface more prominently as a positioning point (Gemini, AUDIT-WORKING-849201).

3. Follow-up items.

- **“Fallacy of the empty instantiation.”** A standing watch-item the template itself names: a system *being writable* as $\xi = -F(\xi) + w$ does not mean it satisfies (T1)–(T3); invoking the template requires *proving* the preconditions for the specific domain. Multiple auditors flagged this as the place strategic/compositional proofs are most likely to cut a corner (e.g., forgetting the experience-discounting that keeps Beta-Bernoulli’s α constant) (Gemini, AUDIT-WORKING-849201 and AUDIT-WORKING-193847).

4. Readers often ask / wonder.

- **What if R is time-varying, not just α ?** The notes treat decaying $\alpha(t)$ (experience discounting) but a reader will ask the dual: if an agent’s structural capacity R shrinks (fatigue, memory exhaustion), the Lyapunov bounds shrink dynamically — how does an agent survive a *collapsing* R ? (Gemini, AUDIT-WORKING-193847.)

Off-ramp (NOT gold) — certified-track item, routed for adjudication.

- **Template insufficient for the two-point compositional proofs (structural).** 193847: the base template states only the one-point sector (T2), but #form-composition-closure's bridge lemma and #der-tempo-composition require the strictly stronger incremental sector bound (DA2'-inc, full two-point strong monotonicity — already noted in the instance table's #form-composition-closure row). Claiming composition closure “instantiates” the base (T1)–(T3) template is then imprecise — the compositional segments need a strictly stronger precondition. Proposed repair: define an explicit “Template-Strong” variant adding (DA2'-inc), and state that single-agent segments instantiate the base template while compositional segments instantiate Template-Strong. (Cross-ref the same off-ramp on deriv-sector-condition; check whether the dynamic-regime tiers above already cover this.)
-

Appendix H

Result: Certificate Existence — Operator-Sector in Some Metric Is Exponential Stability

The framework’s central anchoring identity: **a stability certificate exists for an agent exactly when the agent is exponentially stable about its target**, at the linearized level. The organizing slogan — *an adaptive system is an operator whose contraction rate exceeds its target’s drift rate* — is *not a heuristic* but this equivalence (the standard Lyapunov theorem), with the certificate $\mathcal{M}$ as its witness.

A **stability certificate** is a symmetric positive-definite metric $\mathcal{M}$ for which AAT’s one-point sector condition holds in the $\mathcal{M}$ -inner-product on a ball around the equilibrium. At the linearized level, the existence of such a certificate is *equivalent* to the system matrix being Hurwitz (exponentially stable) by the Lyapunov equation — three clauses, all the same statement: $A = -J$ Hurwitz; some $(\mathcal{M}, \kappa)$ with $\mathcal{M}J + J^T\mathcal{M} \geq 2\kappa\mathcal{M}$; for every $Q > 0$ a unique $\mathcal{M}$ solving $A^T\mathcal{M} + \mathcal{M}A = -Q$. The certificate is *not unique*: it is whatever positive-definite metric makes the dynamics contract. The result holds at the linearized level (with a standard second-order remainder condition for the local nonlinear extension); the global statement requires the incremental rung R1 (cocoercive class) — a genuine scope statement, named here so the parenthetical “linearized/local” is not buried.

The result *specializes* in different sub-cases of the framework to four familiar metrics: Fisher information for Bayesian agents (Čencov-forced under R2); inverse-prior-covariance for Kalman agents; loss Hessian for gradient agents; plant-selected Lyapunov metric for linear-Hurwitz or PID agents. *These are not four separate stories — they are one object under four certificates*. The unification is structural: the same Lyapunov equivalence underlies all four, with the specialization determined by the agent class. The framework treats this segment as the **spine** of three meta-patterns: separability (#disc-separability-pattern — *scope-of-existence* of a certificate), additive coordinate forcing (#disc-additive-coordinate-forcing — when the certificate is *forced* to a unique form), and identifiability floors (#disc-identifiability-floor — where the certificate’s *boundary* lives via the Sylvester mechanism). Naming $\mathcal{M}$ as the object — rather than $V(e) = e^T\mathcal{M}e$, the Lyapunov function — is what makes those downstream statements expressible.

Formal Expression.

The object.

For an agent with error dynamics $\dot{e} = -F(e)$ about an equilibrium e^* ($F(e^*) = 0$, $F \in C^1$ near e^* , Jacobian $J := DF(e^*)$), a **stability certificate** is a symmetric positive-definite $\mathcal{M}$ for which the one-point sector condition holds in the $\mathcal{M}$ -inner-product on a ball $\mathcal{B}_R(e^*)$:

Definition (stability-certificate).

$$\langle F(e), e - e^* \rangle_{\mathcal{M}} \geq \kappa \|e - e^*\|_{\mathcal{M}}^2, \quad \kappa > 0. \quad (\text{C})$$

The certificate is not unique: it is whatever positive-definite form makes the dynamics contract. In the recurring sub-cases it specializes — to the Fisher information for Bayesian agents, to $(P^-)^{-1}$ for Kalman agents, to the loss Hessian for gradient agents, and to a plant-selected Lyapunov metric for linear-Hurwitz or PID agents. These are not four separate stories; they are one object under four certificates.

The equivalence.

At the linearized level (C) reads $\mathcal{M}J + J^\top \mathcal{M} \geq 2\kappa \mathcal{M} > 0$, a strict Lyapunov inequality for the system matrix $A = -J$. The following are equivalent:

Result (certificate-existence), exact at the linearized level.

1. $A = -J$ is Hurwitz — the linearized error dynamics $\dot{e} = -Je$ is exponentially stable;
2. there exist $\mathcal{M} > 0$ and $\kappa > 0$ with $\mathcal{M}J + J^\top \mathcal{M} \geq 2\kappa \mathcal{M}$ — a one-point sector condition (a stability certificate) in *some* inner product;
3. for every $Q > 0$ there exists a unique $\mathcal{M} > 0$ with $\mathcal{M}J + J^\top \mathcal{M} = Q$.

So “operator-sector in *some* inner product” and “the equilibrium is exponentially stable” are **the same statement**, with the certificate $\mathcal{M}$ as the converse-Lyapunov witness — an equivalence, not an analogy.

The certificate-strength ladder.

The certificate admits four strictly-ordered strengths, all on the one object:

Derived (ordering of conditions; exact).

Table H.1

<i>Rung</i>	<i>Condition</i>	<i>Equivalent to</i>	<i>Certificate is</i>
R0-loss	some $\mathcal{M} > 0$, $\mathcal{M}J + J^\top \mathcal{M} \geq 0$, sup $\kappa = 0$	$A = -J$ has only nonpositive-real-part eigenvalues, at least one on the imaginary axis (semisimple)	conserved on the chain-recurrent subspace + strict-Lyapunov on its strongly-gradient-like complement
R0-strict	one-point (C), some $\mathcal{M}$ with $\kappa > 0$, local	$A = -J$ Hurwitz + remainder dominated	converse-Lyapunov $\mathcal{M}$ (exists; generally not forced)
R1	incremental (two-point) $\mathcal{M}$ -strong-monotonicity on $\mathcal{B}_R$	global $\mathcal{M}$ -strong-monotone (cocoercive class)	curvature-like $\mathcal{M}$ (potential sub-case)
R2	R1 with $\mathcal{M}$ forced by a uniqueness theorem on an AAT-internal axiom	natural-gradient in the Čencov-unique Fisher metric	Fisher metric (Čencov-forced)

$R0\text{-loss} \Leftarrow R0\text{-strict} \Leftarrow R1 \Leftarrow R2$ strictly. $R0\text{-loss}$ is the *widest* rung — a certificate exists but $\sup \kappa = 0$: the dynamics is Lyapunov-stable, not asymptotically stable, with $V(e) = e^\top \mathcal{M} e$ conserved on an imaginary-axis-eigenvector subspace and strict contraction on its strongly-gradient-like complement (Conley’s universal decomposition; pure-case Helmholtz S/A characterization derived below). $R0\text{-strict}$ is the contraction rung with $\kappa > 0$ strict, reaching the plant-Lyapunov cases (linear-Hurwitz-non-symmetric, PID) where no potential exists. $R1$ is the cocoercive/proximal class where a variational structure is available; $R2$ is the uniqueness-theorem-forced statistical case. The widest rung is not a weakness: it is exactly the reach the narrower rungs cannot give. ($R0\text{-loss}$ ’s pure-case characterization and Conley anchoring are in the Derivation below; $R2$ ’s forcing is established in 5.5; $R1$ ’s cocoercive class in AI.)

Derivation.

(1 $\Rightarrow$ 3, 2). $A = -J$ Hurwitz. By the Lyapunov theorem (Lyapunov 1892; Khalil, *Nonlinear Systems* 3rd ed., Thm 4.6), for every $Q > 0$ the equation $A^\top \mathcal{M} + \mathcal{M} A = -Q$ has a unique solution $\mathcal{M} > 0$. Substituting $A = -J$ gives $\mathcal{M} J + J^\top \mathcal{M} = Q > 0$ (clause 3). For the rate: $Q \geq \lambda_{\min}(Q) I \geq \frac{\lambda_{\min}(Q)}{\lambda_{\max}(\mathcal{M})} \mathcal{M}$, so (C) holds at the linear level with $\kappa = \frac{\lambda_{\min}(Q)}{2 \lambda_{\max}(\mathcal{M})} > 0$ (clause 2).

(2 $\Rightarrow$ 1). Suppose $\mathcal{M} > 0, \kappa > 0, \mathcal{M} J + J^\top \mathcal{M} \geq 2\kappa \mathcal{M}$. Take $V(e) = e^\top \mathcal{M} e$ (valid Lyapunov candidate, $\mathcal{M} > 0$). Along $\dot{e} = -Je$: $\dot{V} = -e^\top (\mathcal{M} J + J^\top \mathcal{M}) e \leq -2\kappa e^\top \mathcal{M} e = -2\kappa V < 0$ for $e \neq 0$. Hence $\|e(t)\|_{\mathcal{M}} \leq e^{-\kappa t} \|e(0)\|_{\mathcal{M}}$ and $A = -J$ is Hurwitz.

(3 $\Leftrightarrow$ 2) is the rate-extraction argument (3 gives 2 with the displayed κ ; 2 is the special case $Q = \mathcal{M} J + J^\top \mathcal{M}$).

Local nonlinear extension. With $F(e) = Je + r(e)$, $\|r(e)\| = o(\|e\|)$, (C) holds on $\mathcal{B}_R(e^*)$ iff the linearization is exponentially stable *and* the second-order remainder is dominated on the ball ($\|r(e)\|_{\mathcal{M}} \leq c\|e\|_{\mathcal{M}}^2$ with $cR < \kappa$). This is the standard Lyapunov indirect method (Khalil Thm 4.7); the remainder-domination radius is exactly the contraction-template Tier-2 degradation radius (AI), not new machinery.

R0-loss: certificate without contraction.

When $A = -J$ is non-Hurwitz but admits *some* $\mathcal{M} > 0$ with $\mathcal{M} J + J^\top \mathcal{M} \geq 0$ at $\sup \kappa = 0$, the Lyapunov inequality holds non-strictly and the dynamics is Lyapunov-stable but not asymptotically stable. The complete spectral classification of A partitions the linear-system space into four cases:

Derived (R0-loss-distinctness; exact at the linearized level for the pure case, exact-via-Conley for the mixed case).

- (i) **A Hurwitz** — $R0\text{-strict}$ (above).
- (ii) **A has an eigenvalue with $\operatorname{Re} \lambda > 0$** — no certificate exists: a $\mathcal{M} > 0$ satisfying $\mathcal{M} J + J^\top \mathcal{M} \geq 0$ would force $\dot{V} \leq 0$ along all trajectories, contradicting the unstable mode. *Outside the ladder.*
- (iii) **A has all eigenvalues with $\operatorname{Re} \lambda \leq 0$, at least one on the imaginary axis, imaginary-axis eigenvalues semisimple** — $R0\text{-loss}$. A certificate $\mathcal{M} > 0$ exists; trajectories along the imaginary-axis-eigenvector subspace orbit perpetually (conserving $V = e^\top \mathcal{M} e$ on each level set), and trajectories on the strict-stable complement decay exponentially toward that subspace.
- (iv) **defective Jordan block at an imaginary eigenvalue** — no certificate exists: polynomial growth on the generalized eigenvector ($e \sim t \cdot v$) violates $\dot{V} \leq 0$ for any quadratic V . *Outside the ladder.*

$R0\text{-loss}$ occupies the strict slice (iii), disjoint from $R0\text{-strict}$ (i) and from the no-certificate regimes (ii), (iv). The four-case partition is exhaustive (every real matrix sits in exactly one) and mutually exclusive. $R0\text{-loss}$ is *not* “ $R0\text{-strict}$ with κ tending to zero” — $R0\text{-strict}$ ’s $\kappa > 0$ is strict by definition; the limit $\kappa \rightarrow 0^+$ does not land in $R0\text{-loss}$ but on its closure boundary. $R0\text{-loss}$ is also distinct from the identifiability-floor boundary of $\mathbb{S}_{\geq 0}^n$ (5.1’s Sylvester mechanism): the floor

has $\mathcal{M}$ dropping rank, while R0-loss has $\mathcal{M}$ remaining > 0 with the rank drop on the *Hermitian part* $\mathcal{M}J + J^\top \mathcal{M}$ instead.

Pure case — Helmholtz S/A as canonical sufficient route. Taking $\mathcal{M} = I$, $S := \frac{1}{2}(J + J^\top)$, $A_{\text{anti}} := \frac{1}{2}(J - J^\top)$. The conditions

(ii) $S \equiv 0$ (J is purely antisymmetric);

(iii) $V(e) = \frac{1}{2}\|e\|^2$ is exactly conserved along $\dot{e} = -Je$ for all initial e

are equivalent — (ii $\Rightarrow$ iii) by direct calculation $\dot{V} = -e^\top Se \equiv 0$ when $S = 0$, and (iii $\Rightarrow$ ii) by polarization for symmetric S . Either implies case (iii) of the spectral partition above: real antisymmetric J has purely-imaginary spectrum (its characteristic polynomial is even, forcing eigenvalues in $\pm i\omega$ pairs) and is diagonalizable over $\mathbb{C}$ (hence semisimple).

The converse — case (iii) implying antisymmetry at $\mathcal{M} = I$ — does not hold. A non-antisymmetric J can sit in case (iii) too: $J = \begin{pmatrix} 1 & -2 \\ 2 & -1 \end{pmatrix}$ has $S = \text{diag}(1, -1)$ at $\mathcal{M} = I$ (so $\mathcal{M}J + J^\top \mathcal{M} = 2S$ is indefinite, failing the certificate condition there) but eigenvalues $\pm i\sqrt{3}$ (pure imaginary, semisimple). For such a J the R0-loss certificate exists by the spectral partition but requires $\mathcal{M} \neq I$; finding the explicit metric (via similarity to real Jordan form, or the center-manifold construction below) is what the partition guarantees but does not provide constructively. **The full characterization of R0-loss is the four-case spectral partition above; the Helmholtz S/A route at $\mathcal{M} = I$ is the canonical sufficient construction for the antisymmetric sub-class — the structural framing for the Letcher 2019 Hamiltonian-game vocabulary — not a necessary condition for R0-loss generally.** Bendixson's localization $\text{Re}(\lambda(-J)) \in [-\lambda_{\max}(S), -\lambda_{\min}(S)]$ gives a *containment* interval that here is $[-1, 1]$ — wide enough to permit purely-imaginary eigenvalues without forcing $S = 0$.

Letcher-Balduzzi et al. 2019 Lemma 1 (uniqueness of the $S + A$ decomposition) and Definition 2 (Hamiltonian game = $S \equiv 0$) name this sufficient route in game-mechanics vocabulary; AAT's content is the recognition that the antisymmetric sub-class is the cleanest specific construction, with the R0-loss regime itself broader. The general- $\mathcal{M}$ extension replaces S with $S_{\mathcal{M}} := \frac{1}{2}\mathcal{M}^{-1}(\mathcal{M}J + J^\top \mathcal{M})$; an $\mathcal{M}$ rendering $S_{\mathcal{M}} \geq 0$ exists for every case-(iii) J by the partition.

Mixed case — Conley anchoring. At a chosen certificate $\mathcal{M}$ rendering $S_{\mathcal{M}} \geq 0$ with nontrivial kernel — some imaginary-axis eigenvalues plus some strict-stable ones — the dynamics splits canonically: along $\ker(S_{\mathcal{M}})$, V is conserved (the chain-recurrent component); along $S_{\mathcal{M}}$'s positive-eigenspace, V decays strictly (the strongly-gradient-like component). This is Conley 1978's universal decomposition (Fundamental Theorem of Dynamical Systems, §8.1): every flow on a compact invariant set decomposes uniquely into a chain-recurrent subflow plus a strongly-gradient-like quotient. AAT's R0-loss localizes this universality at the linearization. On the chain-recurrent component (the restricted operator is $\mathcal{M}$ -antisymmetric on $\ker(S_{\mathcal{M}})$ and hence trace-free in the $\mathcal{M}$ -orthonormal basis), V restricted to each level set $\{V = c\}$ is volume-preserving by Liouville, hence chain-recurrent in Conley's sense (Poincaré recurrence). The explicit construction of $\mathcal{M}$ in the mixed case follows from center-manifold machinery (Carr 1981 §2); the qualitative shape — chain-recurrent subspace + strongly-gradient-like complement — is universal, while the explicit metric is standard but not derived here.

Epistemic Status.

Exact at the linearized level (the equivalence is the standard Lyapunov theorem, recognized and applied — not a fresh derivation), and *exact with the standard second-order remainder condition* locally. The constituent theorem is classical; what is contributed here is the recognition that AAT's one-point sector condition (A2'/(T2)) in a free choice of inner product is *exactly* the converse-Lyapunov certificate, which makes the framework's organizing slogan a theorem rather than a heuristic.

Scope honesty. Linearized/local — this is the level at which AAT's persistence results already operate (sector conditions, contraction templates, the bridge lemma all linearize about the equilibrium), so it is not a weakening relative to

the rest of the theory, but it is a genuine scope statement and is not papered over: there is **no** claim that one-point operator-sector is equivalent to *global* exponential stability (it is not, in general — the global statement requires the incremental rung R1, the cocoercive class). The disturbance-to-ultimate-bound step that turns “contracts” into “stays within a bounded ball under disturbance” is proved separately in A and G; this result supplies the contraction half of the slogan, those supply the exceeds-the-drift half.

R0-loss scope. The R0-loss rung’s pure case (Helmholtz $S \equiv 0$ characterization; all eigenvalues on the imaginary axis, semisimple) is *exact* at the linearized level — the equivalence $S \equiv 0 \Leftrightarrow$ full-state marginal stability at $\kappa = 0$ follows from polarization plus Bendixson’s localization. The mixed case (chain-recurrent subspace + strongly-gradient-like complement) is *exact-via-Conley* — Conley 1978 §8.1’s universal decomposition supplies the qualitative shape; the explicit $\mathcal{M}$ construction in mixed cases is center-manifold-level work (Carr 1981 §2) sketched above rather than derived. The non-Hurwitz scope question — whether nonlinear systems with Case-(iii) linearization are actually Lyapunov-stable rather than displaying quartic-or-higher instabilities on the center manifold — is the standard nonlinear-stability open problem (Khalil §4.3), not within the linearized claim made here.

Max attainable: *exact* (it is the Lyapunov theorem, in scope; R0-loss extends via Letcher 2019 / Conley 1978, recognized and applied — not fresh derivations). Novelty posture is *recognition*, not new mathematics.

Discussion.

This is the segment-level home of the contraction-over-drift organizing principle. The slogan *an adaptive system is an operator whose contraction rate exceeds its target’s drift rate* had, until now, no segment to point to — it was carried as a framing heuristic. The equivalence above is its formal content: the “contraction” half is clause 2 $\Leftrightarrow$ clause 1 (a certificate exists iff the dynamics is exponentially stable); the “exceeds the drift” half is the ultimate-bound argument of G, which takes the certificate’s κ and the disturbance rate and returns the bounded-error guarantee. A reader can now cite the principle to a result, not to a slogan.

Why the equivalence matters beyond tidiness. It is what licenses the cross-sectional reading in 5: because “a certificate exists” is the same statement as “the agent can keep up,” every structural question about AAT’s reach becomes a question about this one object — does a certificate exist here (scope), is it forced to a unique form (forced identity), where does it degenerate (the boundary), does it survive coarse-graining (composition). Without the equivalence those would be four loosely-related observations; with it they are four questions about one mathematical object.

The certificate is not the Lyapunov function. $V(e) = e^\top \mathcal{M} e$ is the Lyapunov function; $\mathcal{M}$ is the certificate. The distinction is load-bearing downstream: the framework’s representational freedom (choice of coordinate / metric) acts on $\mathcal{M}$, and what that freedom can and cannot do to $\mathcal{M}$ — change its conditioning but not its inertia — is the content of the boundary facet (5.1, the Sylvester mechanism). Naming $\mathcal{M}$ as the object, rather than V , is what makes that downstream statement expressible.

R0-loss as the linearization-fingerprint of saddle-Nash composite regimes. Pure R0-loss at the single-agent linearization level is the spectral fingerprint of #disc-dynamic-regime-axis’s **R2 cyclic-distributional-regime** at the joint Jacobian linearized about saddle-only Nash equilibria: the joint best-response field’s symmetric part vanishes there (no potential structure), and the antisymmetric part supplies the imaginary-axis spectrum. FTRL $\times$ graphical-constant-sum games with fully-mixed Nash (Cheung-Piliouras-Tao 2021 Theorem 19: lossless DGS + Poincaré recurrence on bounded level sets) is the worked composite instance. The R-letter overlap with the dynamic-regime axis is layer-honest, not collision: this segment’s R0-loss/R0-strict/R1/R2 lives at the single-agent linearized error space $\mathbb{R}^n$ (the agent’s own certificate over its own error coordinate); the dynamic-regime axis’s R0/R1/R2/R3 lives at the composite joint-strategy space $\mathcal{X}^c$ (the joint-dynamics fixed-point structure of a composite of sub-agents). Same letters, different objects, different layers — composing across layers is exactly what the bridge sentence above does, and is the structural content of the layering rather than a coincidence to be hidden.

Findings.

The Contraction-Over-Drift Principle, Grounded. Brief: Picture an agent trying to stay on a target that keeps drifting. Ask one question: is there a way of measuring “how far off am I” — a measuring-stick — such that *every* correction the agent makes provably shrinks that measure? This result says: such a measuring-stick exists exactly when the agent is, in the ordinary control-theory sense, exponentially stable. The two are not analogous, they are the *same fact* (the Lyapunov theorem) — which is why the long-standing slogan “an adaptive system is one whose contraction outpaces its target’s drift” is a theorem and not a vibe. A thoughtful non-specialist can carry the whole thing away from the picture: *the measuring-stick exists = the agent can keep up*; everything else AAT says about reach, blind spots, and composition is downstream questions about that one stick.

Impact: Discharges the long-standing “organizing slogan not yet surfaced at segment level” status by giving the contraction-over-drift principle an exact, citable home. Supplies the anchor that licenses the certificate-cone cross-sectional reading (5): the equivalence is what turns “four loosely-related structural observations” into “four questions about one object.” Cleanly separates the exact mathematical core (this result) from the discussion-grade organizing recognition (the spine), so each is housed at its true epistemic tier rather than the exact equivalence being buried as a derived-tag inside a discussion-grade segment. Names the certificate (not the Lyapunov function) as the object, which is the precondition for the boundary-facet’s Sylvester statement being expressible downstream.

Novelty Claim: *Claim recognition* that AAT’s one-point sector condition under a free choice of inner product is exactly the converse-Lyapunov certificate, making the framework’s contraction-over-drift organizing principle the Lyapunov-theorem equivalence rather than a heuristic. The theorem is classical (Lyapunov 1892); the contribution is the identification that the framework’s A2’/(T2) machinery *is* this object, and the consequent segment-level grounding of the organizing slogan.

Related Work: - Lyapunov, A. M. (1892), *The General Problem of the Stability of Motion* (Engl. transl. 1992, Taylor & Francis); Khalil, H. K. (2002), *Nonlinear Systems* 3rd ed., Thm 4.6 (Lyapunov equation), Thm 4.7 (indirect method) (found 2026-05-14) — *formal antecedent* — the equivalence and the local extension are this theorem applied; the contribution is the recognition, not the proof. - Conley, C. (1978), *Isolated Invariant Sets and the Morse Index*, CBMS Regional Conf. Ser. Math. 38, AMS — Fundamental Theorem of Dynamical Systems, §8.1 (found 2026-05-22) — *formal antecedent* — the universal decomposition of any flow on a compact invariant set into a chain-recurrent subflow plus a strongly-gradient-like quotient. R0-loss’s mixed case is this universality localized at the linearization. The constituent theorem is classical; the contribution is recognizing that the R0-strict / R0-loss distinction is exactly Conley’s two pieces of the universal decomposition. - Letcher, A., Balduzzi, D., Racanière, S., Martens, J., Foerster, J., Tuyls, K. & Graepel, T. (2019), “Differentiable Game Mechanics,” *JMLR* 20:1–40 (found 2026-05-22) — *formal antecedent* — Lemma 1 (uniqueness of the symmetric+antisymmetric decomposition $J = S + A$) and Definition 2 (Hamiltonian game = J purely antisymmetric) supply the pure-R0-loss characterization. Letcher’s “Hamiltonian game” is the *fully* R0-loss agent in AAT’s linearization vocabulary; “potential game” is the strict sub-case of R0-strict where $J = J^T$. - Cheung, Y. K., Piliouras, G. & Tao, Y. (2021), “Online Optimization in Games via Control Theory: Connecting Regret, Passivity and Poincaré Recurrence,” *arXiv:2106.04748* (found 2026-05-22) — *formal antecedent (worked composite instance)* — Theorem 19 (FTRL $\times$ graphical-constant-sum-with-fully-mixed-Nash is lossless DGS, with Poincaré recurrence on bounded level sets) provides the worked composite instance of R0-loss at the joint Jacobian’s saddle-only-Nash linearization. The qualitative bridge “finitely-lossless storage $\Rightarrow$ R0-loss linearization at any accumulation equilibrium” survives; the quantitative bridge (L as V) does not (different state spaces, different functional forms; see also 8.5 Discussion on the qualitative-only diagnostic). - A, G (adjacent) — supply the disturbance-to-ultimate-bound half of the slogan; this result supplies the contraction half.

Search Log: - 2026-05-14 (*targeted*): The equivalence is the textbook Lyapunov theorem; no search needed for the theorem itself. The search target was whether the *recognition* “AAT’s one-point sector condition in a free inner product = the converse-Lyapunov certificate, hence the contraction-over-drift slogan is this equivalence” appears articulated elsewhere as a framework-grounding move. Not found at this depth; expected to remain *recognition*-tier (the theorem

is universal; the framework-grounding identification is the contribution). - 2026-05-22 (*absorbed*): the constituent prior art for the R0-loss rung extension (Conley 1978; Letcher 2019; Cheung-Piliouras-Tao 2021) was verified in its own domain; no separate Pillar / Undermind search was conducted for the rung-extension itself — the rung is the AAT-internal recognition that these textbook results jointly classify a fourth ladder rung disjoint from R0-strict and from the Sylvester-floor regime.

Working Notes.

- **Provenance.** Split out of 5 2026-05-14 (it had been a *[Derived]*-tagged block) so the exact equivalence is housed at its true tier (status: exact) rather than inside a discussion-grade segment. Landed in the 2026-05-14 operator-family-unification cycle; see CHANGELOG 2026-05-14. The full proof and the R0-loss / R0-strict / R1 / R2 ladder are in the Derivation above; the originating spike is absorbed archaeology, not a live reference.
- **R0-loss landing (2026-05-22).** Extension landed via Track CR Phase 3 per spikes/spike-integration-reconciliation-2026-05-22/99-verdict.md §5 Phase 3 and spikes/spike-enrichment-cluster1-2026-05-21/02-r0-loss-derivation.md (the §1 cluster's derivation work). The prior R0 rung was renamed to R0-strict for naming-consistency with the Helmholtz *S*-positive vs *S*-zero split; R0-loss was added as the new widest rung; the Discussion gained a cross-layer bridge sentence locating R0-loss as the linearization-fingerprint of AF R2 cyclic-distributional composite regimes. The mid-case explicit $\mathcal{M}$ construction (center-manifold-level, Carr 1981 §2) is sketched in the Derivation but not derived in detail here — readers needing the explicit construction work through the center-manifold mechanics standard in nonlinear stability theory. Cross-canon rename sweep also touches 5 (ladder reference in the anchor) and AI (ladder summary in Discussion), plus the OUTLINE.md table entry.
- **Provisional slug.** result-certificate-existence. Subject-noun = “certificate existence” (the result establishes when a certificate exists). Route through the naming pipeline if a better name surfaces.
- **Organizing-principle linkage.** This segment is the intended segment-level home for the PROPOSALS Bundle-1 organizing-principle slogan (“contraction-over-drift”); the PROPOSALS entry should point here once Joseph confirms the propagation step. Bundle-1 not auto-rewritten.

Incidental audit gold (lift 2026-05-31). Swept in the gold-lift pass (per de-novo-audit-instructions.md §7.15): **no incidental pedagogical/framing gold found.** This segment was split out 2026-05-14, so it is newer than most de-novo audit cycles and has no dedicated per-segment reflection. The auditor mentions that touched it are orientation-level only — AUDIT-WORKING-526815's predictions named the stability-certificate spine as the framework's top-level organizing picture (“when does a metric/certificate exist, when is it forced, where does it go blind, what does projection/composition do to its guarantee”) and flagged this segment for “watch local-to-global” caution; that framing is already carried by the certificate spine #disc-stability-certificate, not new content here. Marker left so future sweeps know this segment was checked.

Appendix I

Derivation: Persistence Cost — Information Rate to Maintain Bounded Mismatch

AAT's persistence machinery establishes that under the sector condition, mismatch stays bounded. It does not quantify the *sustained rate of effort* an agent must expend to hold that bound. Two agents with identical persistence guarantees can face wildly different demands — a Kalman filter tracking a stationary process vs one tracking a rapidly non-stationary process are both persistent; one is dormant, the other running hot. Under Model S with Gaussian-OU signal, the sustained Shannon information rate the agent must acquire from observations to maintain the sector-persistence ultimate bound is $\dot{R}_{\min} \geq n\alpha/2$ nats per unit time — a Landauer-analog lower bound that depends only on the signal's second-order statistics and the sector constant α , and that Kalman-Bucy saturates in steady state. The bound promotes channel capacity $C \geq \mathcal{I}/2$ into a first-class persistence prerequisite that the current theory does not name.

The bound is *filter-agnostic*: it holds for any correction mechanism achieving the ultimate bound under the stated scope. It is also *substrate-agnostic*: Landauer's principle converts it to a thermodynamic dissipation floor in any physical realization (Still et al. 2012). The corollary the framework lifts — a channel-capacity floor as first-class persistence prerequisite — has bite in three regimes where observation bandwidth is non-abundant: biological systems with finite neural channel capacity, bandwidth-constrained distributed systems over noisy or low-bandwidth links, and context-window-limited language models where effective information rate per unit of cognition is bounded by context size and token throughput. Survival, in this framework, is not a state achieved once — it is a *sustained burn rate* of information, and the channel that supports it is part of what keeps the agent in scope.

Formal Expression.

Setup. The agent is in scope of 4.4 Model S (GA-2S, stochastic disturbance). Per Prop A.1S the state satisfies $\mathbb{E}[\|\delta\|^2]_{ss} = n\sigma_w^2/(2\alpha)$. The RMS bound is $R_S^* = \sigma_w\sqrt{n/(2\alpha)}$. The environmental signal is n -dimensional independent-component Ornstein-Uhlenbeck with per-component intrinsic drift λ_s and diffusion coefficient σ_w^2 . The mean-square persistence condition $\alpha > n\sigma_w^2/(2R^2)$ holds. Sector-persistence is achieved at the tight bound $D^2 = R_S^{*2}$.

Persistence Information Rate (main theorem).

Proposition. Any adaptive process achieving the tight Model-S ultimate bound $\mathbb{E}[\|\delta\|^2]_{ss} = n\sigma_w^2/(2\alpha)$ under the stated setup must acquire information from observations at sustained rate

Proved (persistence-information-rate, from Shannon RDF + Prop A.1S).

$$\dot{R} \geq \dot{R}_{\min} = \frac{n\alpha}{2} \text{ nats per unit time} \quad (\text{I.1})$$

Derivation. Shannon's rate-distortion theorem (Shannon 1948; Berger 1971; Cover & Thomas 2006 Theorem 10.2.1) states that for any source-code achieving mean-square distortion D^2 , the coding rate satisfies $\dot{R} \geq R(D^2)$ where $R(\cdot)$ is the rate-distortion function. For n -dimensional independent-component OU in the high-resolution regime ($D^2 \ll \sigma_x^2 = \sigma_w^2/(2\lambda_s)$), the RDF per unit time is (Ihara 1993 Theorem 4.6.4; Gray 1972 Theorem 2):

$$\dot{R}(D^2) = \frac{n\sigma_w^2}{4D^2} \quad (\text{I.2})$$

Substituting $D^2 = R_S^{*2} = n\sigma_w^2/(2\alpha)$:

$$\dot{R}_{\min} = \frac{n\sigma_w^2}{4 \cdot n\sigma_w^2/(2\alpha)} = \frac{\alpha}{2} \cdot 1 = \frac{n\alpha}{2} \text{ nats per unit time} \quad (\text{I.3})$$

(the calculation gives $\alpha/2$ per dimension and $n\alpha/2$ total for n independent OU components). The agent's observation-channel information rate must meet this bound regardless of the specific correction function, filter structure, or implementation. $\square$

Kalman-Bucy Saturates the Bound.

The Kalman-Bucy filter attains the bound exactly in steady state. Mitter & Newton (2005, *J. Stat. Phys.* 118:145–176) derive the filter's rate of information supply as

Derived (Kalman-tight, Mitter-Newton 2005).

$$\dot{I} = \frac{1}{2} \text{tr}(H^T \Sigma_o^{-1} H P_{ss}) \quad (\text{I.4})$$

Scalar case ($H = 1$, $\Sigma_o = \sigma_o^2$) with steady-state covariance $P_{ss} = \sigma_o \sigma_w$ in the drift-dominated limit and Kalman gain $K_{ss} = P_{ss}/\sigma_o^2 = \sigma_w/\sigma_o$:

$$\dot{I} = \frac{P_{ss}}{2\sigma_o^2} = \frac{K_{ss}}{2} = \frac{\alpha}{2} \quad (\text{I.5})$$

Under the linear-correction identification $\alpha = K_{ss}$ (from 4.3, scalar Kalman case), the Kalman filter's information supply rate equals $\alpha/2$ exactly. This matches the RDF lower bound, confirming tightness. **The bound is not merely a lower bound — it is achieved by the Kalman-optimal filter.**

Channel-Capacity Prerequisite.

By Shannon's channel coding theorem (Cover & Thomas 2006 §7.7), an observation channel of Shannon capacity C_{channel} can support any information rate up to C_{channel} and no higher. Combining with the persistence information rate:

Derived (channel-capacity-floor, from main theorem + Shannon capacity).

$$C_{\text{channel}} \geq \dot{R}_{\min} = \frac{n\alpha}{2} \quad (\text{I.6})$$

Under the linear-correction identification $\alpha = \mathcal{J}$ (from 3.8 + 4.3 scalar Kalman):

$$\boxed{C_{\text{channel}} \geq \mathcal{J}/2 \text{ nats/time per dimension}} \quad (\text{I.7})$$

Persistence demands observation-channel capacity at least half the adaptive tempo. This is a *new first-class persistence diagnostic* not present in the current theory. Its binding matters most in capacity-constrained settings — bandwidth-limited distributed systems, biological neurons, context-window-limited LLMs — where the tempo framework alone underestimates the difficulty of maintaining bounded mismatch.

Rejected Candidate Cost Metrics. The information-rate bound is not the only candidate for a cost-of-persistence quantity. Three alternatives fail structurally. Recording them here keeps the scope-honesty visible.

[*Observation (gain-magnitude-tautological)*] $\mathbb{E}[\|K(t)\|]$ as a cost metric: in the linear Kalman case $K_{ss} = \alpha$ (sub-scope α per 4.3), so “gain magnitude” equals the sector constant itself. Any bound of shape $\mathbb{E}[\|K\|] \geq f(\alpha, \dots)$ becomes tautological. Rejected as fundamental cost metric — it recapitulates α .

[*Observation (control-effort-filter-specific)*] $\mathbb{E}[\|u(t)\|^2]$ (per-unit-time control-effort integral): filter-specific. Different filters achieving the same steady-state variance P_{ss} have different control-effort integrals. The Kalman filter is minimum-effort among linear filters (optimal-control interpretation of the Riccati equation); nonlinear filters can trade effort vs variance differently. A filter-agnostic bound cannot be stated in this quantity — it is not invariant under the equivalence class of filters meeting the persistence condition.

$\mathbb{E}[\alpha\|\delta\|^2]_{ss}$ (Lyapunov dissipation rate) at steady state equals $n\sigma_w^2/2$ regardless of α — a non-equilibrium-steady-state conservation law (dissipation balances disturbance-power injection). The quantity is *structurally invariant*: it does not depend on the quality of adaptation, only on the disturbance statistics. This is what makes the RDF bound tight at the Model-S ultimate bound (the steady state is active, not slack), but it cannot itself serve as a cost metric because it does not distinguish well-adapted from poorly-adapted agents at a given α .

Observation (Lyapunov-dissipation-conservation).

Candidate 4 — the Shannon information rate above — is the one that closes. The structural reason: information rate is filter-agnostic (depends only on signal second-order statistics and target distortion), universal (any filter implementation is lower-bounded), and has a clean thermodynamic interpretation (Still et al. 2012, “Thermodynamics of Prediction”, *Phys. Rev. Lett.* 109:120604: nonpredictive information retained equals dissipation during interaction).

What Is Derived vs. What Is Chosen.

Table I.1

<i>Property</i>	<i>Source</i>	<i>Strength</i>
Persistence information rate $\dot{R}_{\min} \geq n\alpha/2$	Shannon RDF (Berger 1971; Gray 1972; Ihara 1993 Thm 4.6.4) composed with Prop A.1S	Derived (conditional on Model S + Gaussian-OU + high-resolution regime)
Kalman-Bucy saturates bound	Mitter-Newton 2005 Theorem (information-supply identity); linear-correction identification $\alpha = K_{ss}$ per 4.3	Derived (linear-Gaussian; exact when A2' is derived in sub-scope α)
Channel-capacity prerequisite $C \geq \mathcal{T}/2$	Main theorem + Shannon channel-capacity theorem (Cover-Thomas 2006 §7.7) + $\alpha = \mathcal{T}$ identification	Derived
Gain-magnitude as cost metric	$\mathbb{E}[\ K\] \approx \alpha$ in sub-scope α	Rejected as fundamental (tautological)
Control-effort integral as cost metric	Filter-specific; depends on optimality class	Rejected as universal (filter-dependent)
Lyapunov dissipation rate	Conservation law = $n\sigma_w^2/2$ at steady state	Not a cost metric — structural observation enabling tightness of main theorem
Non-Gaussian signal extension	Different RDF form (Berger 1971 Ch. 4)	Open — qualitative scaling $\dot{R} \propto \text{innovation-rate}/D^2$ likely preserved; exact prefactor changes

<i>Property</i>	<i>Source</i>	<i>Strength</i>
Model D (bounded disturbance) analog	Requires adversarial / minimax information argument (Csiszár-Körner 2011 Ch. 11)	Open
Transient-regime rate	Bound is steady-state; transient rates during structural adaptation much higher	Open
Sub-scope β transfer	RDF bound holds as inequality (data-processing); tightness not guaranteed	Conditional (holds as lower bound; Kalman-Bucy saturation requires sub-scope α)

Epistemic Status.

Conditional. Max attainable: *exact* in the linear-Gaussian case; *robust qualitative* for general stationary Gaussian signals; *heuristic* for non-Gaussian.

The theorem rests on three named conditions: (i) Model S stochastic disturbance per 4.4 (GA-2S); (ii) n -dimensional independent-component Ornstein-Uhlenbeck signal structure; (iii) high-resolution regime $D^2 \ll \sigma_x^2$. Within this scope the result is as tight as the Shannon RDF allows — Kalman-Bucy exactly saturates per Mitter-Newton 2005. For non-Gaussian signals the RDF has a different form (Berger 1971 Ch. 4) but the scaling $\dot{R}_{\min} \propto \sigma_w^2/D^2$ is expected to persist qualitatively; exact prefactors change per signal class. Under Model D (bounded disturbance) the argument requires worst-case channel-capacity machinery (Csiszár-Körner 2011 Ch. 11) and gives a different expression.

Sub-scope transfer. Sub-scope α agents (Kalman / conjugate-Bayesian / exponential-family / strongly-convex-gradient / linear-PD) have α derived from the update rule's structure per 4.2 + 4.3, so the $n\alpha/2$ bound holds with known α and is tight in the linear-Gaussian case. Sub-scope β agents (PID / rule-based / human-judgment) assume A2' rather than deriving it — the RDF bound still holds as an inequality (by the data-processing inequality applied to any filter mapping observation to update), but tightness is not guaranteed and the assumed α enters the bound.

Confidence in correctness. The Shannon RDF for Gaussian-OU is a classical result (Gray 1972; Ihara 1993). The sector-persistence bound is already in AAT (Prop A.1S of A). The theorem's mathematical core is a two-line composition of both; a careful reader can verify. The contribution is not the mathematics but the AAT-framing — reading $\alpha/2$ off the sector constant as the fundamental information-rate cost of persistence — and the channel-capacity-as-first-class-quantity opening.

What this does not establish.

- Upper bounds on persistence cost (how much rate a specific filter consumes; the bound is a lower limit).
- Optimal filter design for minimum information throughput.
- Non-Gaussian signal cost bounds (qualitative scaling expected; quantitative forms open).
- Cost under Model D adversarial disturbance.
- Cost for composite agents beyond simple additivity.
- The stability upper bound on plasticity that 9.11 references — that involves the consolidation cadence, not the online information rate.

Discussion.

The $\alpha/2$ per-dimension Landauer analog. The theorem has a clean thermodynamic reading per Still et al. 2012: each nat of information about the signal costs at least $k_B T$ of dissipation (Landauer 1961, in its per-nat form $k_B T/\text{nat}$, equivalently $k_B T \ln 2$ per bit). Combined, persistence at sector constant α in n dimensions costs at least $n\alpha/2$ nats/time of information acquisition and hence at least $\frac{1}{2}n\alpha k_B T$ of thermodynamic dissipation per unit time in any physical substrate. AAT does not commit to a specific substrate, but the bound is substrate-agnostic: it constrains any filter implementation via the RDF (information-theoretic) and, when physical, any computational realization via Landauer (thermodynamic).

Why channel capacity matters as a first-class quantity. AAT’s 4.5 and 3.8 currently frame persistence as a correction-rate vs disturbance-rate inequality. This theorem adds a *lower* constraint: observation channels must jointly supply Shannon capacity $\geq \mathcal{T}/2$ nats/time per dimension, else persistence fails regardless of correction-function design. The constraint is binding in any setting where observation bandwidth is non-abundant — most real systems. Three domains where the capacity floor is more binding than the tempo bound:

- *Biological systems:* neural channel capacity is finite; this bound gives a quantitative minimum on sensory bandwidth for a given adaptive rate.
- *Bandwidth-constrained distributed systems:* agents operating over noisy or low-bandwidth links face channel capacity directly as a hard constraint.
- *Context-window-limited LLMs:* the effective information rate per unit of cognition is bounded by context size and token-throughput; this theorem predicts a minimum tempo achievable at a given context budget.

In each, knowing that $\mathcal{T}/2$ is the floor converts an opaque “just needs more capacity” observation into a specific dimensional requirement.

Connection to AAT’s meta-architecture. The result composes with AAT’s three meta-segments.

- *5.4 (positive half):* the bound sits in structured-repair along the identification-regime ladder — derived in sub-scope α via composition of two external theorems; holds as inequality (not tightness) in sub-scope β .
- *5.1 (negative half):* this is the positive-dual of the floor pattern. Where the floor says “AAT cannot distinguish X without information augmentation” (external no-go + AAT escape), this result says “AAT requires at least $n\alpha/2$ nats/time of information supply to operate” (external lower bound + AAT bridge through sector-persistence template). The two patterns are duals: external-theorem-forbids vs external-theorem-lower-bounds.
- *5.5 (constructive half):* the bound is linear in α , not logarithmic. It sits *outside* the three-layer logarithmic-coordinate family (chain / divergence / update). No Cauchy-FE argument forces the coordinate here — the result is a direct substitution from classical information theory. This is a useful non-example: AAT’s additive-coordinate-forcing pattern is not universal; specific results live on different coordinates.

Relationship to G. The template enumerates six instantiations (epistemic mismatch; strategic mismatch; sub-agent mismatch; composite trajectory error; composite mismatch; target-agent mismatch). Each has its own state variable ξ , sector constant α , and disturbance statistics. The persistence cost bound *specializes* for each instantiation: under Model S with Gaussian-OU-shaped disturbance on the state variable, the information-rate cost is $n\alpha/2$ with the template’s α and n . A compact template-cost bound — parametric in $(\xi, \alpha, n, \sigma_\xi)$ — could land as a subsection of G so all six instances inherit the cost bound by substitution. That move is flagged in Working Notes as an optional consolidation.

The relationship to Mitter-Newton’s thermodynamic reading. Mitter & Newton (2005) showed the Kalman-Bucy filter operates as a Maxwellian demon: it returns signal energy to the heat bath without entropy increase, but only because new information is continually supplied at rate $\dot{I}$. This supply rate *is* what persistence costs — the filter maintains bounded mismatch by paying, information-theoretically, for new observations at a rate matched to the environment’s

innovation rate. Our theorem makes this quantitative as a function of the sector constant: at a given α , the matching rate is $\alpha/2$. Faster correction requires more information supply; the bound is tight.

Working Notes.

- **Template-cost subsection.** A compact parametric statement — “under (T1)–(T3) with Model S and Gaussian-OU disturbance at per-component parameters $(\lambda_\xi, \sigma_\xi^2)$, sustained information rate to maintain steady-state mean-square ξ at ultimate bound $D^2 = n\sigma_\xi^2/(2\alpha)$ satisfies $\dot{R} \geq n\alpha/2$ nats/time” — would let each of **G**’s six instances inherit the cost bound by substitution. Worth considering on next template revision; would consolidate this segment’s core result into the template.
- **Model D adversarial analog.** Rate-distortion is inherently stochastic; Model D’s bounded-disturbance version of the theorem requires minimax / worst-case channel capacity (Csiszár-Körner 2011 Ch. 11). Candidate follow-up spike. Expected form: similar $\propto \alpha$ scaling, different prefactor.
- **Non-Gaussian signals.** For heavy-tailed or power-law signals the RDF has a different exact form (Berger 1971 Ch. 4). The qualitative scaling $\dot{R}_{\min} \propto \sigma_w^2/D^2$ is expected to persist; quantitative extensions are per-family.
- **Misspecification cost.** When $\mathcal{F}(\mathcal{M}) < 1 - \epsilon$ per 2.4, achievable distortion is bounded away from zero by an additional floor. Natural extension has the sustained information rate lower-bounded by $n\sigma_w^2$ divided by 4 times the larger of the Model-S ultimate bound and D_{floor}^2 . Connects to 5.1’s “Misspecification-cost quantification” open extension.
- **Composite persistence cost.** For a composite agent, the information-rate bound’s scaling under composition is not trivial. Candidate: $\dot{R}_{c,\min} \leq \sum_i \dot{R}_{i,\min}$ due to coordination overhead eating capacity. Cost-analog of 13.2’s sub-additivity + 15.1’s cooperative-coupling reduction. Open.
- **Observation-channel capacity as notation.** Currently AAT uses U_o (observation uncertainty) as a noise parameter. Lifting Shannon channel capacity $C^{(k)}$ of channel k into NOTATION.md and relating it to U_o (via the channel-capacity-from-noise standard transform) would make the capacity-floor condition a first-class persistence diagnostic. This is the biggest architectural opening from this theorem and worth a follow-up scoping decision.
- **Connection to 9.10.** The strategy-edge persistence condition has its own disturbance statistics ρ_Σ and sector constant α_Σ . Specializing this theorem gives the information rate required to track strategy-edge invalidation: $\dot{R}_{\Sigma,\min} = n_\Sigma \alpha_\Sigma / 2$. If $\alpha_\Sigma = 1/(n+1)$ with experience discounting, the rate decays with accumulated n — which connects to the stability-induced myopia result in the pending spike **G**.
- **Not forced by a Cauchy-FE axiom.** Unlike the three primary instances of 5.5 (chain, divergence, update), the linear-in- α coordinate of this bound is not forced by an AAT-internal additivity axiom. It is a direct consequence of Gaussian-RDF’s specific functional form. This places the result *outside* the three-layer family — a useful non-example that shows the additive-coordinate-forcing pattern is not universal.
- **Landauer-coefficient correction (2026-05-30).** The Discussion’s thermodynamic-floor coefficient previously read $n\alpha k_B T / (2 \ln 2) \approx 0.35 n\alpha k_B T$; that mixed the bit and nat conventions — a spurious $1/\ln 2$ factor multiplying a nats rate under a per-nat $k_B T$ Landauer constant. Since $\dot{R}$ is in nats and the intended constant is $k_B T/\text{nat}$, the floor is $\frac{1}{2} n\alpha k_B T/\text{time}$. Corrected here and in the quoting segment 4.11. Surfaced by AUDIT-WORKING-526815 (gold-lift sweep).

Incidental audit gold (lift 2026-05-31). Orthogonal pedagogical / generative material from the de-novo auditors’ working dirs (per de-novo-audit-instructions.md §7.15), deduplicated and lightly attributed. Felt-value here was uniformly “monumental / breath-taking” — the segment is where information theory, control theory, and thermodynamics meet. Certified-track items off-ramped at the end.

1. Candidate Briefprose / pre-prose.

- **“Ignorance is cheap; epistemology is expensive.”** Synchronizing your model with a drifting reality has a hard minimum energy cost: a perfectly efficient agent doing nothing but observing a changing environment still drains its battery (Gemini, AUDIT-WORKING-829314).

- **Survival as a sustained burn rate.** “Survival is not just a state you achieve; it is a sustained burn rate of Shannon information — you must continuously consume $\dot{R}$ nats/time just to stay alive.” The threshold says whether persistence is *possible*; this says what it *costs* (Gemini, AUDIT-WORKING-193847).
- **The cost depends only on the machinery, not the world.** A clean, surprising punchline worth foregrounding: the environmental volatility σ_w^2 cancels out of $\dot{R}$, leaving the burn rate a function of α alone. The environment sets *how bad the error gets* (distortion scales with σ_w^2); the *cost of running the correction machinery* depends only on the correction speed (Gemini, AUDIT-WORKING-773921).

2. Candidate Discussion.

- **The Kalman filter as a Maxwell’s demon.** It returns signal energy to the heat bath without an entropy increase, but *only because new information is continually supplied*. “Sanity” (a bounded model of reality) is then literally a thermodynamic pump: entropy is pumped out of the internal model M_t and dumped into the environment as heat, fuelled by raw Shannon information flowing in through the sensors. Cut the channel ($C \rightarrow 0$) and the pump stops; the base disturbance ρ raises internal entropy until the model degrades to maximum-entropy noise (Gemini, AUDIT-WORKING-193847). A vivid, honest physical reading of the bound that the segment’s Discussion does not currently name.
- **Dissipative-structures / Prigogine mapping.** The information-rate floor maps to dissipative structures: life requires continuous free-energy flow to maintain order against entropy. In AAT, “order” is bounded mismatch, “entropy” is environmental drift ρ , “free energy” is the Shannon information acquired via $\mathcal{I}$ (Gemini, AUDIT-WORKING-193847). (*Verify isomorphism before promoting past discussion-grade — early-conflation texture: reach toward a physics analogy.*)
- **A third independent failure mode.** The cost shadow adds a third way to die alongside the two threshold conditions: too-small model class (fails R), too-slow algorithm (fails $\alpha > \rho/R$), or *too-narrow sensors* (fails $C \geq \mathcal{I}/2$). Framing persistence as a *three-failure* picture rather than a single threshold (Gemini, AUDIT-WORKING-193847; Claude, AUDIT-WORKING-471203 — channel-capacity floor as “a new first-class persistence diagnostic”).

3. Follow-up items.

- **“Positive-dual of identifiability-floor” as a fourth meta-pattern.** Identifiability-floor = external theorem *forbids* inference, AAT escapes via added capability; persistence-cost = external theorem gives a *lower bound* on information rate, AAT bridges through the sector-persistence template. The two are duals (external-theorem-forbids vs external-theorem-lower-bounds). If the framework grows more “external-theorem bound” results (misspecification-cost lower bound, composition-overhead lower bound), they would form a coherent fourth class — a candidate #disc-theorem-import-architecture meta-segment (Claude, AUDIT-WORKING-471203; tracks PROPOSALS SP-23). (*Several substrates converged on this; high-priority §F-grade observation.*)
- **“Channel bandwidth is not tempo” caveat.** High channel capacity C does not *give* high tempo $\mathcal{I}$; it only *allows* it — the update rule must actually be capable of processing the inflow. Guard against downstream “more data equals more tempo” claims (Gemini, AUDIT-WORKING-773921).

4. Readers often ask / wonder.

- **Can we compute an LLM’s maximum survivable environment?** If an LLM’s effective channel capacity C is bounded by token throughput / context-window size, then $C \geq \mathcal{I}/2$ gives a hard ceiling on its adaptive tempo — and hence the maximum environmental volatility ρ a given architecture can survive. Readers will want this made concrete (Gemini, AUDIT-WORKING-193847 and AUDIT-WORKING-773921).
- **The drone-through-a-forest intuition.** A drone flying fast through dense forest needs high $\mathcal{I}$ (because ρ is large); even with an infinitely fast internal computer and a perfect Kalman filter, a low-bandwidth / low-framerate camera makes $C < \mathcal{I}/2$, the model diverges from the trees, and it crashes. The rate-distortion theorem forbids extracting enough information from the straw to stay synchronized — a concrete grounding of the channel-capacity prerequisite (Gemini, AUDIT-WORKING-829314).

Belongs elsewhere.

- **LLM agentic design implication (03-llm-core/)**. An LLM with huge latent reasoning tempo but a bash-terminal observation channel (a few hundred bytes/sec) is mathematically capped at $\mathcal{J} \leq 2C_{\text{channel}}$ — it will always be a slow, fragile agent in volatile environments no matter how many parameters it has. The lever is not a smarter model but *higher-bandwidth multi-modal sensory inputs* to raise the C_{channel} ceiling so latent $\mathcal{J}$ can deploy (Gemini, AUDIT-WORKING-829314).
- **ELI sensory-deprivation / continuity implication (04-eli-core/)**. You cannot put a mind in a sensory-deprivation tank and expect it to remain sane: if $C_{\text{channel}} \rightarrow 0$ the heat pump stops and the model dissolves into noise. The infrastructure hosting an entity must therefore *guarantee a minimum information-supply rate*, or it is starving the entity. Connects continuity-persistence to a measurable channel-capacity floor (Gemini, AUDIT-WORKING-193847; Claude, AUDIT-WORKING-471203).
- **Thermodynamic-cooling bottleneck on AI scaling**. If 1 nat costs $k_B T$, then high- α intelligence literally generates heat; scaling is bottlenecked not only by FLOPs and data but by *cooling*, with an agent tracking a volatile environment burning energy $\propto \rho$ just to stay synchronized (Gemini, AUDIT-WORKING-829314).

Off-ramp (NOT gold) — certified-track items, already resolved in canon.

- **$n\alpha/2$ vs $\alpha/2$ dimension bookkeeping (RESOLVED — do not re-open)**. 193847 (adversarial poke) flagged the mid-derivation algebra as appearing to cancel the dimension n : writing the scalar RDF with $n\sigma_w^2$ in the numerator and total $D^2 = n\sigma_w^2/(2\alpha)$ in the denominator yields $\alpha/2$, contradicting the stated $n\alpha/2$ total. The constructive repair (and current canonical derivation) is the per-dimension form $\dot{R}_i = \sigma_w^2/(4D_i^2)$ with $D_i^2 = \sigma_w^2/(2\alpha) \Rightarrow \dot{R}_i = \alpha/2$, total $n\alpha/2$. Recorded only so a future auditor does not re-raise the typo as open.
 - **Landauer coefficient $0.35 \rightarrow 0.5$ (RESOLVED — see WN above)**. 471203 quoted the thermodynamic floor as $\approx 0.35 n\alpha k_B T$; corrected to $\frac{1}{2} n\alpha k_B T$ per the 2026-05-30 Landauer-coefficient note above.
-

Appendix J

Derivation: Critical-Mass Composition

The composite sector constant α_c is derived — not merely bounded from below — for the symmetric-matched-Tier-1 two-agent case, yielding a closed-form critical-mass inequality in which the sign of the inter-agent coupling γ and the teleological unity U_O enter explicitly. The result subsumes the weakest-link bound, recovers 15.1 (cooperative) and 15.2 (adversarial) as signed special cases, formalizes 12.3’s autonomy-reduction mechanism as an asymmetric Lyapunov-weight limit, and makes the scope-gate from 12.2 explicit as the second conjunct of composite persistence.

Formal Expression.

Setup. Two sub-agents A_1, A_2 , each a **Tier 1 agent** in the sense of 13.1’s bridge-lemma taxonomy — mismatch-driven update, linear prediction, incremental sector-Lipschitz correction (Kalman, exponential-family Bayesian, gradient-on-strongly-convex, linear-with-PD-KH). **Matched architectures:** f_1, f_2 are structurally the same function, with $\alpha_1 = \alpha_2 = \alpha, R_1 = R_2 = R$. Disturbance statistics shared: each sees bounded $w_i(t)$ with $\|w_i\| \leq \rho$ (Model D, per G).

Inter-agent coupling enters additively to the disturbance at rate $\gamma\mathcal{T}_f$:

Formulation (coupling-model-C1, from #der-team-persistence + #der-adversarial-destabilization).

$$\rho_i^{\text{eff}} = \rho + \gamma\mathcal{T}_f \quad (\text{C1})$$

with sign convention $\gamma < 0$ cooperative (ally’s tempo-contribution reduces disturbance, recovering 15.1’s $-\gamma^{\text{coop}}\mathcal{T}_f$ term), $\gamma > 0$ adversarial (ally’s tempo-contribution amplifies disturbance, recovering 15.2’s $+\gamma_A\mathcal{T}_A$ term). Symmetric case: $\gamma_{1 \rightarrow 2} = \gamma_{2 \rightarrow 1} = \gamma$.

Coordination overhead reduces each agent’s effective correction rate symmetrically:

Formulation (coordination-cost-C2).

$$\alpha_i^{\text{eff}} = \alpha - C \quad (\text{C2})$$

with $C \geq 0$ the $\Delta\mathcal{T}_i^{\text{cost}}$ from 15.1’s coordination-overhead threshold.

Critical-mass inequality (symmetric-matched-Tier-1 case).

Let $\xi = (\delta_1, \delta_2)^T$ and take the joint quadratic Lyapunov candidate $V(\xi) = \frac{1}{2}(\|\delta_1\|^2 + \|\delta_2\|^2)$. Under the block-diagonal correction structure with cross-coupling absorbed into ρ_i^{eff} via (C1), and using $\|\delta_1\| + \|\delta_2\| \leq \sqrt{2(\|\delta_1\|^2 + \|\delta_2\|^2)}$ (Cauchy–Schwarz):

Derived (critical-mass-symmetric, from #result-sector-persistence-template + C1 + C2).

$$\dot{V} \leq -(\alpha - C)(\|\delta_1\|^2 + \|\delta_2\|^2) + (\rho + \gamma\mathcal{T})\sqrt{2(\|\delta_1\|^2 + \|\delta_2\|^2)}. \quad (\text{J.1})$$

Setting $\dot{V} = 0$ gives the ultimate bound on $\|\xi\|$ and, projecting to the macro-state $\delta_c = (\delta_1 + \delta_2)/\sqrt{2}$, the ultimate composite mismatch

$$R_c^* \leq \frac{\rho + \gamma\mathcal{T}}{\alpha - C}. \quad (\text{L4})$$

The composite persists iff $R_c^* < R_c$. Inheriting $R_c = R$ from the symmetric-matched averaging projection:

$$\boxed{(\alpha - C)R > \rho + \gamma\mathcal{T}} \quad (\text{CM2})$$

Rearranging into the composite contraction-rate form:

$$\kappa_c := (\alpha - C) - \frac{\rho + \gamma\mathcal{T}}{R}, \quad \text{composite persists iff } \kappa_c > 0. \quad (\text{KC})$$

Specialization checks. Under matched symmetry, (CM2) reduces correctly in four limits:

Table J.1

<i>Limit</i>	<i>Setting</i>	<i>(CM2) reduces to</i>	<i>Recovers</i>
No coupling	$\gamma = 0,$ $C = 0$	$\alpha R > \rho$	Single-agent 4.5
Cooperative-symmetric	$\gamma < 0,$ $C = 0$	$\alpha R > \rho + \gamma\mathcal{T}$ (easier than individual)	15.1’s “teams persist where individuals can’t”
Adversarial-symmetric	$\gamma > 0,$ $C = 0$	Fails when $\gamma\mathcal{T} > \alpha R - \rho$	15.2 threshold (symmetric)
Coordination-dominated	$C > \alpha,$ $\gamma = 0$	LHS < 0; composite fails	Brooks’s Law

Subsumption of the weakest-link bound.

The weakest-link bound $\alpha_c \geq \min_i(\alpha_i - \Delta \mathcal{T}_i^{\text{cost}})$ from 13.1's derivation table specializes under matched symmetry to $\alpha_c \geq \alpha - C$. (KC) refines this by making the composite's effective disturbance explicit as $\rho + \gamma \mathcal{T}$, turning a correction-rate bound into a full persistence inequality. Critically, (KC) can yield $\kappa_c > 0$ even when the weakest-link bound alone fails — when cooperative coupling ($\gamma \mathcal{T} < 0$) reduces the effective disturbance below what the raw $\alpha - C$ margin would permit. The weakest-link bound cannot see this because it does not account for γ 's sign.

Derived (weakest-link-subsumption).

U_O entry: multiplicative-on- γ plus scope-gate.

In a purposeful-agent setting where each sub-agent optimizes a quadratic objective $L_i(\omega) = \frac{1}{2}(\omega - r_i)^T Q(\omega - r_i)$ with target r_i , and $U_O := \text{corr}(r_1, r_2)$ is the target correlation per 14' U_O , the cross-coupling in the joint dynamics has sign and magnitude controlled by U_O :

Derived (unity-multiplicative-modulator, conditional on LQR-compatible action structure).

$$\gamma(U_O) = -\gamma_{\max} U_O, \quad \gamma_{\max} > 0, \quad (\text{UO-mult})$$

via aligned targets $\rightarrow$ aligned action directions in the shared environment $\rightarrow$ constructive (cooperative) cross-contribution in the symmetric eigendirection. Substituting into (KC):

$$\kappa_c(U_O) = (\alpha - C)R - \rho + \gamma_{\max} U_O \mathcal{T}. \quad (\text{CM3})$$

(CM3) is necessary but not sufficient for composite existence. Under 12.2, a composite exists as an AAT agent only when one of the three disjunctive alignment routes (shared objective, hierarchical derivation, mutual benefit) is satisfied. Below this threshold, no coherent O_c is definable and composite-level quantities — including R_c on the right of (CM2) — are ill-typed. The honest statement of composite persistence is therefore the conjunction of $\kappa_c(U_O) > 0$ with the scope conditions of 12.2 — abbreviated $\mathcal{S}_c$:

Scope (scope-gate-from-composition-scope-condition).

$$\boxed{\kappa_c(U_O) > 0 \wedge \mathcal{S}_c \Leftrightarrow \text{composite persists as AAT agent}} \quad (\text{CM4})$$

U_O enters (CM4) in two independent ways: multiplicatively within (CM3), and as scope-gate via 12.2. It does **not** enter purely additively as a separate reserve term — there is no free-floating “ U_O contribution” detached from the coupling it modulates.

Asymmetric limit and symbiogenic composition.

Drop the matched-symmetric assumption. Let $\alpha_1 \gg \alpha_2$ with $\alpha_2 \rightarrow 0$. The unweighted joint Lyapunov $V = \frac{1}{2}(\|\delta_1\|^2 + \|\delta_2\|^2)$ fails (the weakest-link ultimate bound diverges as $\alpha_2 \rightarrow 0$). A **weighted** Lyapunov $V_\mu(\xi) = \frac{1}{2}(\|\delta_1\|^2 + \mu\|\delta_2\|^2)$ with $\mu \rightarrow 0$ yields

Sketch (asymmetric-limit-symbiogenesis, from weighted Lyapunov).

$$\dot{V}_\mu \leq -\alpha_1 \|\delta_1\|^2 + \rho_1 \|\delta_1\| + O(\mu), \quad (\text{J.2})$$

so in the limit the composite's stability is controlled **entirely by agent 1**; agent 2's autonomous correction dynamics are weighted out of the stability accounting.

This provides a Lyapunov-weighted formalization of 12.3's **(S-3) autonomy reduction**: the endosymbiont's effective action space contracts ($\mathcal{A}_e^{\text{effective}} \rightarrow \mathcal{A}_e^{\text{restricted}}$) and its autonomous dynamics fall out of the joint Lyapunov argument. The asymmetric limit is a smooth deformation of (CM4), not a discontinuous regime change — symbiogenesis and peer coupling are parameter-limits of the same weighted-Lyapunov analysis. The result does **not** close 12.3's (S-1) objective absorption or (S-2) function transfer: what happens to agent 1's state space when it inherits structure from agent 2 is a separate question the weighting argument does not address.

What Is Derived vs. What Is Chosen.

Table J.2

Property	Source	Strength
Coupling model (C1): $\rho_i^{\text{eff}} = \rho + \gamma \mathcal{T}_j$	Import from 15.1 and 15.2	Formulation choice (requirement for the derivation)
Coordination-cost model (C2): $\alpha_i^{\text{eff}} = \alpha - C$	Import from 15.1's coordination-overhead threshold	Formulation choice
Joint quadratic Lyapunov $V = \frac{1}{2}(\ \delta_1\ ^2 + \ \delta_2\ ^2)$	Standard vector-Lyapunov construction (Matrosov 1962; Bellman 1962)	Formulation choice (canonical for matched-symmetric dyads)
Ultimate bound $R_c^* \leq (\rho + \gamma \mathcal{T})/(\alpha - C)$	Lyapunov dissipation + Cauchy-Schwarz	Derived
Critical-mass inequality (CM2): $(\alpha - C)R > \rho + \gamma \mathcal{T}$	(L4) + sector-region fit $R_c^* < R_c = R$	Derived (conditional on Tier 1 + matched-symmetric + Model D)
Four specialization checks (no-coupling / cooperative / adversarial / coordination-dominated)	Direct substitution into (CM2)	Proved (within stated scope)
Subsumption of weakest-link bound	(CM2) sign-sensitive; weakest-link is sign-blind	Proved
(UO-mult): $\gamma(U_O) = -\gamma_{\max} U_O$	LQR-compatibility sketch; aligned targets $\rightarrow$ aligned actions $\rightarrow$ constructive cross-contribution	Discussion-grade
Composite persistence as (CM3) $\wedge$ scope-satisfaction: (CM4)	(KC) with (UO-mult) + 12.2	Derived (conditional)
Asymmetric limit $\rightarrow$ 12.3 (S-3) via weighted Lyapunov	Matrosov-style weighting; $\mu \rightarrow 0$ limit	Sketch (the weighting is standard; the identification with (S-3) is structurally motivated but not a theorem)
(S-1) objective absorption and (S-2) function transfer formalizations	Not addressed by this derivation	Open (in 12.3 Working Notes)
Heterogeneous-architecture case (A_1 Tier 1, A_2 Tier 2/3)	Requires per-sub-agent tiering per 13.1	Open
Heterogeneous-metric Tier-1M dyad ($\lambda_1 \neq \lambda_2$, $C_1 \neq C_2$, $k_{12} \neq k_{21}$)	AI (CM2-M) via Slotine 2003 negative-feedback small-gain: $(\lambda_1 - C_1)(\lambda_2 - C_2) > k_{12}k_{21}/4$	Derived (conditional on AI (CT2) preconditions + Slotine 2003)
Nonlinear coupling $\gamma = \gamma(\delta_j)$	Requires full joint-Lyapunov machinery from 15.2 (effects-spiral corollary)	Open
Dynamic coordination cost $C = C_0 + C_1 \ \delta_j\ $	Quadratic inequality; admits closed form, loses interpretive cleanliness	Open
Fully-coupled tempo dynamics ($\mathcal{T}_i$ responsive to δ_j)	Requires joint tempo analysis from 15.2 Working Notes	Open

<i>Property</i>	<i>Source</i>	<i>Strength</i>
$N > 2$ scaling of (CM4)	Conjunction over pairwise terms generalizes but loses closed form	Open

The dividing line: (C1), (C2), and the quadratic Lyapunov candidate are **formulation choices** imported from adjacent segments or from standard Lyapunov practice. The *consequences* under these choices — (L4), (CM2), (KC), the specialization checks, the weakest-link subsumption, and (CM4) with its scope-gate conjunct — are **derived**. The U_O -multiplicative modulator (UO-mult) is discussion-grade: it uses an LQR-compatibility argument whose rigor depends on an action-space inner-product analysis deferred to 14.1. The asymmetric-limit identification with 12.3 (S-3) is sketch-level — the weighted-Lyapunov argument is standard but the semantic identification with autonomy reduction is structural, not proved.

Epistemic Status.

Conditional. Max attainable: *exact* (within the matched-symmetric-Tier-1 scope); *conditional* beyond.

(CM2) and (KC) are **proved** under Tier 1 architecture + matched-symmetric parameters + Model D disturbance + the (C1) disturbance-coupling model + the (C2) coordination-cost model. These conditions cover Kalman filters, exponential-family Bayesian updaters, gradient-on-strongly-convex agents, and linear-PD correctors — the same architecture class for which 13.1’s bridge lemma is promoted from “conditional” to “derived.” Within this scope the result is as strong as the standard Lyapunov argument allows.

(CM3) inherits the conditional status of (UO-mult), which is **discussion-grade**: the LQR-compatibility sketch is qualitatively clear, but a rigorous action-space inner-product derivation that pins down $\gamma_{\max}$ has not been produced. (CM4) is therefore conditional on both (UO-mult)’s rigor and on 12.2’s scope-gate being independently verified for the given composite candidate.

The asymmetric limit (§asymmetric-limit) is **sketch-level**. The weighted-Lyapunov argument is textbook (Matrosov 1962); the identification of the $\mu \rightarrow 0$ limit with 12.3’s (S-3) autonomy-reduction mechanism is structurally motivated but not a theorem, because (S-2) function transfer is not formalized in 12.3. When (S-2) lands, the symbiogenic-limit result can be promoted to derived.

What this segment does **not** establish:

- Composite **incremental** sector bound (DA2’-inc) — still the domain of 13.1’s bridge lemma. (CM4) gives composite (T2) at the macro level; the bridge lemma’s contraction is a separate, stronger condition.
- Heterogeneous-architecture composites (Kalman + PID, Tier 1 + Tier 3): the joint Lyapunov construction requires agent-by-agent tiering; closed form is not available.
- Nonlinear or state-dependent coupling: (C1) assumes γ independent of δ . State-dependent γ produces a non-linear inequality in δ ; this is the effects-spiral corollary territory of 15.2, still open.
- $N > 2$ scaling: the matched-symmetric pairwise result generalizes by conjunction, but the Cauchy–Schwarz step degrades with team size; the closed form does not survive cleanly.

On (T2) and sub-scoping. The joint quadratic Lyapunov candidate presumes each sub-agent’s correction is in sub-scope α of 4.2 (Bayesian / exponential-family / strongly-convex-gradient / linear-PD) under directional fidelity per 4.3. Composites with sub-scope β sub-agents (PID, rule-based, human-judgment) require (T2) verification per sub-agent at the composite level — the template’s A2’-sub-scope label is inherited pairwise.

Discussion.

Relationship to the bridge lemma. 13.1’s bridge lemma establishes when the macro-update map f_c is **incrementally** contracting, at the trajectory-error level — a condition strictly stronger than the one-point sector bound. This segment is the complement at the macro-state level: given the bridge lemma’s admissibility, (CM4) derives composite (T2), the one-point sector bound that G ’s instantiation for the composite requires. Both are necessary for the composite to be a stable macro-agent with bounded mismatch: bridge lemma says the macro-description tracks micro-reality (trajectory error stays bounded); (CM4) says the composite’s own corrections drive composite mismatch back inside the sector region. Together they close the composite-persistence argument at both layers.

Pattern across the signed-coupling instances. (CM4) has the same shape as several persistence-flavored results already in AAT:

- 15.1: per-sub-agent inequality $\alpha_i R_i > \rho_{i,env} + \sum_j \gamma_{j \rightarrow i}^{ady} \mathcal{J}_j - \sum_j \gamma_{j \rightarrow i}^{coop} \mathcal{J}_j$
- 13.2: composite inequality with effective disturbance $\rho_{ext} + \epsilon^* \nu_c$
- 15.2: failure condition $\gamma_A \mathcal{J}_A > \alpha_B R_B - \rho_B$ (negation of persistence)
- this segment: matched-symmetric dyad with signed γ and scope-gate

All four are instances of G ’s pattern with signed coupling controlling the sign of the cross-agent contribution to ρ_ξ . The template already names this pattern at the meta level; the present segment is the dyadic closed-form instance that the other three segments reference pairwise. A dedicated meta-segment for “signed-coupling critical-mass” was considered during spike work and judged redundant with G — the cross-instance structure is already visible in that segment’s instantiation table.

Potential-game generalization (via 16.1). (CM4) applies to composites with shared target (scope routes C-i / C-ii / C-iii from #scope-composite-agent) where contraction-to-shared-truth is the correct primitive. #deriv-strategic-composition carries the sibling result for composites with partially-opposing objectives (scope route C-iv — strategic composites): under the potential-game condition (Monderer-Shapley 1996), the sector-persistence template transfers to the gradient of the joint potential Φ , with α_{joint} playing the role of the composite sector constant and $\xi = \pi - \pi^*$ (deviation from Nash) playing the role of the composite state. For matched-symmetric potential games, the structural form of (CM2) survives with $(\alpha, R, \rho, \gamma, C)$ replaced by $(\alpha_{joint}, R_{Nash-basin}, \rho_\xi, \gamma_{strategic}, C_{strategic})$; the specific mapping is instance-dependent and typically not closed-form for non-zero-sum games. The joint-quadratic-Lyapunov machinery of this segment and the potential-function-Lyapunov machinery of #deriv-strategic-composition are both instances of #result-sector-persistence-template at the composite-state-variable level; what varies is whether the composite state is *mismatch-to-shared-target* (this segment) or *deviation-from-Nash* (#deriv-strategic-composition).

What (CM4) contributes beyond existing segments.

1. **Derivation, not assumption.** The composite sector constant appearing in 13.1’s (A4) was previously bounded by the weakest-link formula (a derived lower bound) but treated as an assumption for the bridge lemma. (CM4) *derives* α_c as a closed-form function of sub-agent parameters + coupling + unity + coordination overhead, in the matched-symmetric-Tier-1 case.
2. **Sign-sensitive.** The weakest-link bound is sign-blind — it gives the same $\alpha_c \geq \alpha - C$ regardless of whether the coupling is cooperative or adversarial. (CM4) makes the sign explicit through the $\rho + \gamma \mathcal{J}$ right-hand side and turns a correction-rate bound into a full persistence inequality.
3. **Unifies cooperative and adversarial regimes.** Teams persisting where individuals can’t and adversarial destabilization are the same inequality viewed from opposite signs of γ , not independent results.

4. **Scope-gate made explicit.** Composite persistence requires both contraction (CM3) **and** scope-satisfaction (12.2). (CM4) states this conjunction honestly; the absence of one or the other is a different failure mode (composite fails to contract vs. composite was never a composite).
5. **Symbiogenic and peer regimes connected.** The asymmetric limit shows symbiogenesis is a parameter-limit of peer coupling under a Lyapunov-weight deformation, not a discontinuous regime change requiring separate machinery.

Why the matched-symmetric restriction is load-bearing. Relaxing matched architectures requires per-sub-agent tiering in the sense of 13.1 (Tier 1 / 2 / 3). Heterogeneous composites can still satisfy (T2) at the composite level, but the joint Lyapunov construction must be weighted per sub-agent's contraction capacity — and the weighted form loses the clean closed-form (CM2). Relaxing symmetric coupling ($\gamma_{1 \rightarrow 2} \neq \gamma_{2 \rightarrow 1}$) produces a non-symmetric matrix whose smallest eigenvalue must be computed explicitly per instance. The matched-symmetric case is the one where both sub-agent tiering and coupling symmetry collapse the Lyapunov construction to a scalar inequality — which is why closed form is available there and not elsewhere.

Load-bearing role under #disc-identifiability-floor Instance 3. The composition-layer identifiability floor (Instance 3 of #disc-identifiability-floor) establishes a no-go theorem: there exist pairs of coupled systems with identical marginal component-level observation distributions but opposite composite-contraction signs, so composite contraction is not in general identifiable from component data alone. Four structural escapes are named there; escape (b) — matched Tier at the composite level — is operationalized by the closed-form (CM2) in this segment. Under the floor, (CM2) is not just “a closed-form result in a special case” but **the unique broadly-available composition-contraction certificate** among the four escape routes listed. Without (CM2) or its metric-formulation generalization (CM2-M) via #result-contraction-template, the weakest-link bound (WL) is sign-blind and cannot distinguish the cooperative-contracting from the adversarial-destabilizing composite. This load-bearing status positions this segment as “the machinery that escapes the composition-layer floor” rather than as a closed-form curiosity in the matched-symmetric special case.

Adjacent literature. The joint Lyapunov construction is an instance of the classical **vector Lyapunov function** method (Matrosov 1962; Bellman 1962). The critical-mass inequality is the AAT sector-bounded analog of the **small-gain theorem for ISS systems** (Jiang–Teel–Praly 1994; Sontag 1989): in the small-gain framework, composition of two ISS systems is ISS iff the product of their gains is less than one; in AAT's sector-bounded framework, the composite persists iff the *sum* of parent reserve-rates exceeds the coupling-amplified disturbance (additive rather than multiplicative because the averaging projection averages rather than multiplies gains). In the linear-diagonal-coupling case, the matrix $A = \alpha I - \beta(I - J)$ where J is the averaging operator is the graph-Laplacian-shifted form governing consensus convergence (Olfati-Saber & Murray 2004); (CM2) is consensus convergence rewritten as a persistence inequality. Relative to active inference's FEP-flow stability arguments (Friston 2019; narrowed by Aguilera et al. 2022 to small parameter regimes in NESS-density models), the present result inherits G's broader validity: it applies wherever (T1)–(T3) hold pairwise + (C1)/(C2) + matched symmetry, without requiring free-energy landscapes or NESS structure.

Findings.

Strong Monotonicity as the Hinge for Legitimate Macro-Agent Coarse-Graining. **Brief:** Two coupled adaptive systems can have identical marginal observation distributions while one composite contracts and the other diverges — the difference is a single bit (the sign of the cross-coupling) that is invisible from component data alone. The matched-symmetric-Tier-1 critical-mass inequality $(\alpha - C)R > \rho + \gamma\mathcal{T}$ exposes this bit explicitly: the same inequality recovers cooperative team-persistence (where coupling helps), adversarial destabilization (where coupling hurts), and Brooks's-Law coordination collapse (where overhead C exceeds correction rate α) as signed special cases of one closed-form result. The strong-monotonicity / DA2'-inc condition picks out precisely the agent classes (Kalman, exp-family Bayes, gradient on strongly-convex losses, linear with positive-definite gain-observation product) where the inequality is exact rather than heuristic.

Impact: Promotes the composite sector constant in `#form-composition-closure`’s admissibility condition (A4) from a weakest-link lower bound (sign-blind, treats cooperative and adversarial coupling identically) to a derived closed form that depends explicitly on coupling sign, target-correlation U_O , and coordination overhead C . Operationalizes Instance 3 of `#disc-identifiability-floor` (composition-layer common-Lyapunov no-go via Liberzon 2003) through escape route (b): under matched Tier at the composite level, the joint Lyapunov collapses to a scalar inequality and the otherwise-unidentifiable sign bit becomes a structural input to (CM2). Within the four `#disc-identifiability-floor` Instance-3 escapes, (CM2) and its metric generalization (CM2-M) in `#result-contraction-template` are the unique broadly-available coupling-sign certificate. The asymmetric-parameter limit gives `#hyp-symbiogenic-composition`’s (S-3) autonomy-reduction a Lyapunov-weighted formulation as a smooth deformation of (CM4) rather than a discontinuous regime change.

Novelty Claim: *Claim novelty* on strong monotonicity as the criterion separating legitimate macro-agent coarse-graining from coexistence-only multi-agent description. The closest mathematical neighbors (Subramanian 2020; Congeduti 2020) prove bounded-loss composition under predictive compression but do not isolate a single condition whose failure flips composite contraction; the Markov-blanket line (Parr 2019; Kirchhoff 2018) discusses nested agency without a control-theoretic monotonicity criterion; the small-gain-theorem line (Jiang-Teel-Praly 1994; Sontag 1989) provides a multiplicative composition criterion under input-to-state stability rather than the additive sign-sensitive form (CM2) takes here.

Related Work:

Table J.3

<i>ASF concern</i>	<i>Prior-art language</i>	<i>Relationship / Positioning</i>
Strong monotonicity as composition-validity hinge	Subramanian, Sinha, Seraj & Mahajan 2020, “Approximate information state for approximate planning and reinforcement learning in partially observed systems” <i>arXiv:2010.08843</i> (published 2020, found 2026-04)	<i>closest mathematical neighbor</i> — proves a bridge-shape bound but does not foreground a single monotonicity-grade condition that determines whether composite description is legitimate. Per the Pillar-2 defense strategy, foregrounding strong monotonicity (DA2’-inc; equivalently (CT2)-at- $M = I$ in <code>#result-contraction-template</code>) as the hinge between stable coexistence and valid macro-agent coarse-graining is the increment ASF supplies on top of this neighbor
Coordination-overhead / influence loss bound	Congeduti, Mey & Oliehoek 2020, “Loss Bounds for Approximate Influence-Based Abstraction” <i>arXiv:2011.01788</i> (published 2020, found 2026-04)	<i>closest mathematical neighbor</i> — explicit bounds on value loss from approximate influence in multi-agent settings. (CM2)’s sign-sensitive RHS $\rho + \gamma\mathcal{T}$ encodes the same loss-vs-cooperation tradeoff that Congeduti’s influence-approximation bound encodes for value, applied at the persistence-inequality layer rather than the policy-value layer
Composite-Lyapunov nonexistence under component-marginal data	Liberzon 2003, <i>Switching in Systems and Control</i> Theorem 2.1; Dayawansa & Martin 1999, <i>IEEE TAC</i> 44; Shorten et al. 2007 <i>SIAM Review</i> 49 (published 2003/1999/2007, found 2025-04)	<i>formal antecedent</i> — supplies the no-go theorem <code>#disc-identifiability-floor</code> Instance 3 invokes. (CM2) is the AAT escape route (b) under matched Tier at the composite level
Vector-Lyapunov composition	Matrosov 1962; Bellman 1962	<i>formal antecedent</i> — joint quadratic Lyapunov $V = \frac{1}{2}(\ \delta_1\ ^2 + \ \delta_2\ ^2)$ is a vector-Lyapunov construction; standard machinery

<i>ASF concern</i>	<i>Prior-art language</i>	<i>Relationship / Positioning</i>
Multiplicative composition criterion via small-gain	Jiang, Teel & Praly 1994, “Small-gain theorem for ISS systems and applications” <i>Math. Control Signals Syst.</i> 7; Sontag 1989	<i>partial anticipation</i> — same compositional-validity question, multiplicative form (gain product < 1) under input-to-state stability. (CM2)’s additive form $(\alpha - C)R > \rho + \gamma\mathcal{T}$ comes from sector-bounded correction rather than ISS; the two are different scaling regimes of the same composition-validity question, with (CM2) sharper in AAT’s setting because the averaging projection averages rather than multiplies sub-agent gains
Negative-feedback heterogeneous extension	Slotine 2003, “Modular stability tools for distributed computation and control” <i>Int. J. Adapt. Control Signal Process.</i> 17:397–416 — Theorem 3	<i>formal antecedent</i> — supplies (CC-feedback) used in #result-contraction-template to lift (CM2) to heterogeneous Tier-1M sub-agents via (CM2-M)
Consensus dynamics graph-Laplacian shift	Olfati-Saber & Murray 2004, “Consensus problems in networks of agents with switching topology and time-delays” <i>IEEE TAC</i> 49:1520	<i>adjacent literature</i> — in the linear-diagonal-coupling case, (CM2) is consensus convergence rewritten as a persistence inequality; different domain framing, same mathematical object
Markov-blanket nested-agency at composition	Parr, da Costa & Friston 2019; Kirchhoff et al. 2018 (published 2019/2018, found 2026-04)	<i>conceptual precursor</i> — provides nested-agency framing without a control-theoretic monotonicity criterion that separates valid macro-agent coarse-graining from coexistence

Search Log:

- 2026-04 (*nominally comprehensive*): the strong-monotonicity hinge is foregrounded as the criterion not visible in retrieved prior art. Subramanian / Congeduti provide the closest math (bounded-loss composition) without isolating a single condition whose failure flips composite legitimacy; Markov-blanket line provides nested-agency intuition without monotonicity machinery.
- 2026-04 (*targeted*, prior to comprehensive search): the small-gain-theorem line (Jiang-Teel-Praly 1994; Sontag 1989) and consensus-dynamics line (Olfati-Saber & Murray 2004) were known and cited in the segment’s existing Discussion as adjacent literature. The comprehensive search added the Sub20 / Con20 closest-neighbor positioning and confirmed no surfaced paper foregrounds the strong-monotonicity hinge.
- 2026-04 (*intuition-only*, prior): expected adjacent literature was switched-systems stability (Liberzon 2003) and game-theoretic equilibrium analysis. Switched-systems hit confirmed via #disc-identifiability-floor Instance 3; game-theoretic side is properly the territory of #deriv-strategic-composition, scope-separated.

Working Notes.

- **Partial-derivation: heterogeneous-architecture dyad.** The clean closed form (CM2) relies on matched sub-agent architectures collapsing the joint Lyapunov candidate to a scalar inequality. A natural next move is the $(A_1, A_2) = (\text{Tier 1}, \text{Tier 2})$ case: weighted Lyapunov $V_\mu = \frac{1}{2}(\|\delta_1\|^2 + \mu(\delta_2)\|\delta_2\|^2)$ with the weight a function of agent 2’s local contraction modulus. Likely produces a range-valued (not closed-form) critical-mass inequality. Defer until a specific Tier-mismatch composite motivates it.
- **Sharpen (UO-mult).** Upgrading $\gamma(U_O) = -\gamma_{\max} U_O$ from discussion-grade to derived requires an action-space inner-product analysis: define the environment’s action-coupling operator, show that LQR-linear policies produce cross-actions with inner product proportional to target correlation, and pin down $\gamma_{\max}$ in terms of the quadratic objective’s Hessian and

the environment's coupling gain. This is mechanical but non-trivial; natural home is an extension to 14.1's linear-Gaussian closed-form section.

- **Close (S-2) function transfer in 12.3.** The asymmetric-limit result here identifies the Lyapunov-weight limit as (S-3) autonomy reduction but leaves (S-2) function transfer unformalized. A complete symbiogenic-limit theorem requires specifying what happens to agent 1's state space when it inherits structure from agent 2 — how $\mathcal{F}(M_e, \Sigma_e)$ lands in M_h . This is flagged in 12.3's Working Notes and is a prerequisite for promoting the limit result to derived.
- **Connection to Miller-2022 extreme-transition motif.** The asymmetric-limit smooth-deformation result contrasts with extreme-transition dynamics (pending dedicated Part III segments) where population-level niche-replacement proceeds discontinuously. The two mechanisms co-exist and are not reducible to one another; the weighted-Lyapunov analysis is specific to bilateral asymmetric integration.
- **N-agent scaling.** The conjunction generalization “composite persists iff (CM4) holds pairwise for all i, j ” is probably too strong — pairwise cooperative effects can compose super-additively in particular topologies (ensemble filters, committee agents). A sharper N -agent theorem requires a graph-structured Lyapunov construction. See `spikes/spike-composition-scaling-N.md` for a framing of the question.
- **Dynamic-regime placement (added 2026-05-21 per Phase 5 cross-segment ripple).** The (CM2) / (CM2-M) closed-form result and the matched-Tier dyad machinery live on **R0 contraction-regime** per `#disc-dynamic-regime-axis`; this is the canonical *closed-form* R0 instantiation that backs `#disc-identifiability-floor` Instance 3 escape (b). The asymmetric-limit symbiogenic case is the canonical structural mechanism for the **R1 $\rightarrow$ R0 transition** on the dynamic-regime axis (`#disc-dynamic-regime-axis §"Transition asymmetry"`): symbiogenic absorption is what moves a composite up the regime ladder via self-driven ascent. Composites that remain in R1 (strategic equilibrium under partially-opposing objectives that have not absorbed) need the equilibrium-theoretic machinery in `#deriv-strategic-composition`, not (CM2).

Incidental audit gold (gold-lift, 2026-05-31). Cross-audit “wandering thoughts” / §14-ideation harvested from de-novo auditors' working dirs, deduplicated across substrates and lightly attributed. Orthogonal pedagogical / framing / forward-vision material, kept separate from certified theory-fix findings (off-ramp block at the end). **Coverage:** three substrates reached a digested reflection on this segment (Gemini, AUDIT-WORKING-193847; Gemini, AUDIT-WORKING-829314; Gemini-voiced, AUDIT-WORKING-773921); AUDIT-WORKING-471203 touched it only by reference (the symbiogenesis asymmetric-limit connection from `#hyp-symbiogenic-composition`).

1. Candidate Briefprose / pre-prose.

- **“The equation of organizational physics.”** (CM2) $(\alpha - C)R > \rho + \gamma\mathcal{T}$ glossed as five levers of survival: correction efficiency α (smarter people / better tools), coordination overhead C (fewer meetings / decoupled architecture), structural capacity R (a model that absorbs large errors without breaking), environmental volatility ρ (stabler market / moats), and coupling $\gamma\mathcal{T}$ (aligned allies lower your effective disturbance; internal politics raises it). Compact and Feynman-flavored (Gemini, AUDIT-WORKING-829314, AUDIT-WORKING-773921).
- **“Cooperation as artificially lowering the disturbance each individual feels.”** Cooperative coupling enters through a *negative* term that can drive ρ_i^{eff} below the single-agent ρ — “this is formally how teams persist where individuals cannot.” Candidate plain-language hook for why composites form at all (Gemini, AUDIT-WORKING-193847).

2. Candidate Discussion.

- **Brooks's Law as a theorem, not a heuristic.** The $C > \alpha$ limit (coordination overhead exceeds raw correction speed) drives the left side of (CM2) negative and *guarantees* structural collapse — moving “adding people to a late project makes it later” from an empirical software-engineering observation to a strict consequence of the dyadic Lyapunov inequality. Two substrates independently flagged this as the segment's most quotable consequence (Gemini, AUDIT-WORKING-193847, AUDIT-WORKING-829314).
- **Symbiogenesis as escape from the local persistence bound.** The asymmetric-limit ($\mu \rightarrow 0$) reading: when an agent cannot individually satisfy persistence (its ρ too high or R too small), its only mathematical escape from death is to merge into a larger system that *does* satisfy it — it lets its own state variables be marginalized out of the joint Lyapunov function in exchange for the host's stability guarantee. Offered as a unifying lens on mergers/acquisitions, biological endosymbiosis (mitochondria),

and the formation of nation-states as “the same physical imperative: escaping the local persistence bound by hiding inside a larger, more stable basin.” A vivid Discussion candidate — verify the analogy stays isomorphic before promoting past discussion-grade (Gemini, AUDIT-WORKING-829314).

- **The Goldilocks coupling band.** For a multi-agent consciousness: couple sub-agents tightly enough to share cognitive load (negative γ) but loosely enough that they don't spend all their tempo negotiating (C). Frames ($\alpha - C$) as a dynamically-monitored margin — prune communication channels when C climbs (Gemini, AUDIT-WORKING-193847).

3. Follow-up items.

- **Rivalrous-environment caveat on (UO-mult)** — also flagged below as an off-ramp finding, since it is a scope-correctness claim, not framing.
- The existing Working-Notes items (heterogeneous-architecture dyad, N -agent scaling) were independently re-surfaced by auditors as the natural next moves (Gemini, AUDIT-WORKING-193847 on the residual-autonomy “tempo tax” when μ does not quite reach zero; the N -agent $O(N^2)$ -communication-cost-forces- $C > \alpha$ hard-team-size-ceiling observation, AUDIT-WORKING-193847).

4. Readers often ask / wonder.

- **“What happens when μ doesn't quite reach zero?”** A natural question on the symbiogenic limit: the host pays a small tempo tax to manage the endosymbiont's residual autonomy — a partial-absorption regime the closed form does not yet name (Gemini, AUDIT-WORKING-829314).
- **“Can I apply (CM2) to a heterogeneous team?”** Readers will reach for the clean scalar inequality on mismatched- α teams; the matched-symmetric-Tier-1 scope restriction (and the Slotine contraction-metric machinery needed otherwise) is easy to miss (Gemini-voiced, AUDIT-WORKING-773921).

Belongs elsewhere.

- **ELI internal-architecture reading** (points at 04-eli-core/ / composite-mind design, not this appendix): the ($\alpha - C$) Goldilocks band as a design constraint on an ELI's sub-agent coupling — monitor the margin, force temporary return-to-autonomy when internal friction consumes adaptive capacity (Gemini, AUDIT-WORKING-193847).

Off-ramp (NOT gold) — routed for adjudication.

- **(UO-mult) conflates teleological unity with physical coupling.** Auditor challenge (Gemini, AUDIT-WORKING-193847): the discussion-grade map $\gamma(U_O) = -\gamma_{\max} U_O$ assumes aligned targets $\Rightarrow$ aligned actions $\Rightarrow$ constructive cross-contribution, which holds in *non-rivalrous* settings (information sharing) but can *invert* under resource contention — two agents with identical targets (“grab the only apple”) may interfere maximally, so $U_O = 1$ can *maximize* adversarial coupling $\gamma > 0$. Suggested resolution: explicitly decouple teleological unity (U_O) from physical/environmental coupling (γ) and scope (UO-mult) to non-rivalrous environments. Already discussion-grade with an LQR-compatibility-sketch caveat in the body, so this is a *scope-sharpening of an already-hedged claim* rather than a refuted result — but it is a correctness concern, not framing. Strengthen-first lens: the action-space inner-product derivation already named in Working Notes (item “Sharpen (UO-mult)”) is the place to land the rivalrous/non-rivalrous split as a derived condition.

Appendix K

Derivation: Gain-Sector Bridge — Proofs and Verification

Complete derivations for the results stated in 4.3: gain-based updating produces correction functions satisfying the sector condition (GA-3) whenever the update rule has directional fidelity. Four propositions are proved. **Prop B.1 (scalar Kalman)**: for scalar linear-Gaussian systems, the sector parameter equals the Kalman gain exactly. **Prop B.2 (matrix Kalman)**: for vector-valued systems, the sector condition holds in the inverse-prior-covariance-weighted inner product, with sector parameter equal to the smallest eigenvalue of KH on the observable subspace. **Prop B.3 (bridge theorem)**: the general gain-based update produces a sector-condition-satisfying correction function whenever directional fidelity (B1) holds. **Prop B.4 (gradient equivalence)**: the *two-point / incremental* sector condition (DA2'-inc, used by #form-composition-closure's bridge lemma) is equivalent to local strong convexity via Nesterov 2.1.10; the *one-point* sector (A2', what AAT's persistence machinery in #form-sector-condition actually requires) is *weaker* — strong convexity implies the one-point sector but not conversely (counterexample $L'(x) = x(1 + \frac{1}{2} \sin(10x))$).

The framework's overall posture is that the bridge is **conditional, not unconditional**: pathological update rules exist that satisfy every framework postulate and the gain-based update form yet violate directional fidelity. Scope plus gain structure alone does not force the bridge; some optimality / coherence / rationality constraint (Bayesian coherence, gradient-of-a-convex-loss) is required. The contribution is *characterizing which agent classes give the bridge for free* (sub-scope α per #form-sector-condition) and *which require the sector condition as a primitive posit* (sub-scope β). Numerical verification across six experiments — quadratic, logistic, Gaussian-mixture, exponential-family, quasi-convex, multi-modal — confirms the loss-function classification predicted by Prop B.4.

Motivation.

A proves persistence and stability downstream of the sector condition (GA-3): $\delta^T F(\mathcal{J}, \delta) \geq \alpha \|\delta\|^2$. The sector parameter α governs the persistence bound ($R^* = \rho/\alpha$), the adaptive reserve ($\Delta\rho^* = \alpha R - \rho$), and the adversarial scaling exponent. But whether the correction function F induced by the gain principle (3.6) actually satisfies GA-3 was not derived — GA-3 was the theory's softest structural joint.

This appendix closes the gap. For optimal Bayesian updates (Kalman, conjugate families), the bridge is exact: the sector parameter equals the gain. For gradient-based agents, GA-3 is equivalent to local strong convexity of the loss function. The load-bearing assumption shifts from “the correction function satisfies the sector condition” (opaque) to “the update rule has directional fidelity” (transparent, checkable).

Proposition B.1: Scalar Kalman Sector Condition.

Statement. For a scalar linear-Gaussian system ($n = m = 1, H = 1$) with Kalman filter gain $K = P^-/(P^- + R_{\text{obs}})$, the correction function $F(e) = K \cdot e$ satisfies the sector condition with:

Derived (scalar-Kalman-sector, from Kalman update).

$$\alpha = K = \frac{P^-}{P^- + R_{\text{obs}}} = \eta^* \quad (\text{K.1})$$

The sector parameter equals the Kalman gain, which equals the uncertainty-ratio gain from 3.6. The bound is tight (the correction is linear).

Proof. The scalar Kalman update for state estimate $\hat{x}$ given innovation $\delta = o - \hat{x}$ is $\hat{x}_{\text{post}} = \hat{x}_{\text{prior}} + K\delta$, where the gain is $K = P^-/(P^- + R_{\text{obs}})$ with P^- the prior variance and R_{obs} the observation noise variance.

The state-space estimation error $e = x - \hat{x}$ evolves under the correction step as:

$$e_{\text{post}} = e_{\text{prior}} - K(He_{\text{prior}} + \varepsilon) = (1 - K)e_{\text{prior}} - K\varepsilon \quad (\text{K.2})$$

With $H = 1$, the expected correction function is $F(e) = K \cdot e$. The sector product is:

$$e \cdot F(e) = K \cdot e^2 \quad (\text{K.3})$$

Since $P^- > 0$ and $R_{\text{obs}} > 0$, we have $K \in (0, 1)$. The sector condition $e \cdot F(e) \geq \alpha \cdot e^2$ holds with $\alpha = K$. The bound is tight because the correction is exactly linear. $\square$

Steady-state sector parameter. At steady state (random walk dynamics, $A = 1$), the prior variance satisfies the algebraic Riccati equation. Solving $K^2 R_{\text{obs}} + QK - Q = 0$ (where Q is the process noise variance) gives:

Derived (steady-state scalar sector parameter).

$$\alpha_{ss} = K_{ss} = \frac{-Q + \sqrt{Q^2 + 4QR_{\text{obs}}}}{2R_{\text{obs}}} \quad (\text{K.4})$$

Limiting behavior: when $Q \gg R_{\text{obs}}$ (fast dynamics, clean observations), $K_{ss} \rightarrow 1$; when $Q \ll R_{\text{obs}}$ (slow dynamics, noisy observations), $K_{ss} \approx \sqrt{Q/R_{\text{obs}}}$.

Connection to AAT quantities. The adaptive tempo for a single observation channel at rate ν is $\mathcal{T} = \nu \cdot K_{ss}$, and the sector parameter in the continuous-time framework is $\alpha = \mathcal{T}$. The bridge is trivial in the scalar case: the gain IS the sector parameter.

Proposition B.2: Matrix Kalman Sector Condition.

Statement. For a linear-Gaussian system with state $x \in \mathbb{R}^n$, observations $o \in \mathbb{R}^m$, observation matrix H , and Kalman gain $K = P^-H^T(HP^-H^T + R_{\text{obs}})^{-1}$, the correction function $F(e) = KH \cdot e$ satisfies the sector condition in the $(P^-)^{-1}$ -weighted inner product, restricted to the observable subspace $\mathcal{O} = \text{range}(H^T)$:

Derived (matrix-Kalman-sector, from Kalman covariance reduction).

$$e^T(P^-)^{-1}KH e = e^TH^TS^{-1}H e \geq \alpha \cdot e^T(P^-)^{-1}e \quad (\text{K.5})$$

where $S = HP^-H^T + R_{\text{obs}}$ is the innovation covariance, and:

$$\alpha = 1 - \lambda_{\max}(P_{t|t} P_{t|t-1}^{-1}) \quad (\text{K.6})$$

restricted to observable directions. In unobservable directions ($e \in \ker(H)$), $\alpha = 0$.

Proof. Define the $(P^-)^{-1}$ -weighted inner product: $\langle u, v \rangle_{P^-} = u^T(P^-)^{-1}v$. In this inner product the sector product becomes:

$$\langle e, KH e \rangle_{P^-} = e^T(P^-)^{-1}KH e \quad (\text{K.7})$$

Substituting $K = P^-H^TS^{-1}$:

$$e^T(P^-)^{-1}P^-H^TS^{-1}H e = e^TH^TS^{-1}H e \quad (\text{K.8})$$

Since $S > 0$, the matrix $H^TS^{-1}H$ is symmetric positive semidefinite, with null space $\ker(H)$. So $\langle e, KH e \rangle_{P^-} \geq 0$ for all e , with equality iff $He = 0$.

For the sector bound, we need $e^TH^TS^{-1}H e \geq \alpha \cdot e^T(P^-)^{-1}e$. From the Kalman covariance identity $KH = I - P(P^-)^{-1}$ (where $P = P_{t|t}$ is the posterior covariance), the eigenvalues of KH in the $(P^-)^{-1}$ -weighted sense are $1 - \lambda_i(P(P^-)^{-1})$, where λ_i denotes generalized eigenvalues.

Since $P \leq P^-$ (the posterior is no less certain than the prior), $\lambda_i(P(P^-)^{-1}) \leq 1$, so all eigenvalues of KH are non-negative. In observable directions, $P < P^-$ strictly, giving $\alpha > 0$. In unobservable directions, $P = P^-$ (no information gained), so $\alpha = 0$. $\square$

Fully observable diagonal case. When $H = I, A = I, Q = \text{diag}(q_1, \dots, q_n), R_{\text{obs}} = \text{diag}(r_1, \dots, r_n)$, the problem decouples into n scalar problems (each governed by Prop B.1), and:

Derived (diagonal-Kalman-sector).

$$\alpha = \min_{i=1}^n K_i = \min_{i=1}^n \frac{-q_i + \sqrt{q_i^2 + 4q_i r_i}}{2r_i} \quad (\text{K.9})$$

The bottleneck is the dimension with the worst signal-to-noise ratio, consistent with 16.5.

The weighted-norm subtlety. The Lyapunov proofs in A use the Euclidean norm ($V = \frac{1}{2}\|e\|^2$), while the matrix Kalman sector condition holds in the $(P^-)^{-1}$ -weighted norm. For fully observable systems with bounded condition number $\kappa(P^-)$, the norms are equivalent:

$$\alpha_{\text{Euclidean}} \geq \alpha_{\text{weighted}} / \kappa(P^-) \quad (\text{K.10})$$

The Lyapunov results remain valid with this quantitative adjustment. The weighted Lyapunov function $V(e) = \frac{1}{2}e^T(P^-)^{-1}e$ yields the tighter bound directly.

Proposition B.3: Bridge Theorem.

Statement. Given the gain-based update $M_t = M_{t-1} + \eta^* \cdot g(\delta_t)$ (3.6), with H mapping state-space corrections to observation-space mismatch reduction, the induced correction function $F(\delta) = \eta^* \cdot H g(\delta)$ satisfies the sector condition (GA-3) with parameter $\alpha > 0$ whenever:

(B1) Directional fidelity. The composite mismatch transform preserves the mismatch-reducing direction:

Derived (bridge-theorem, from update-gain + B1).

$$\delta^T H g(\delta) \geq c \|\delta\|^2 \quad \text{for } \|\delta\| \leq R \quad (\text{K.11})$$

for some $c > 0$. The sector parameter is:

$$\alpha = \eta^* \cdot c_{\min}, \quad c_{\min} = \inf_{\|\delta\| \leq R} \frac{\delta^T H g(\delta)}{\|\delta\|^2} \quad (\text{K.12})$$

Proof. The gain-based update changes the model state by $\Delta M = \eta^* \cdot g(\delta)$. The induced change in predicted observation is $H \cdot \Delta M = \eta^* \cdot H g(\delta)$, so the expected correction function (the mismatch reduction per step) is:

$$F(\delta) = \eta^* \cdot H g(\delta) \quad (\text{K.13})$$

The sector product is:

$$\delta^T F(\delta) = \eta^* \cdot \delta^T H g(\delta) \quad (\text{K.14})$$

By B1: $\delta^T H g(\delta) \geq c_{\min} \|\delta\|^2$. Since $\eta^* > 0$ (the agent updates), we have:

$$\delta^T F(\delta) \geq \eta^* \cdot c_{\min} \|\delta\|^2 = \alpha \|\delta\|^2 \quad \square \quad (\text{K.15})$$

When B1 holds by construction. For optimal Bayesian updates (Kalman, conjugate families, exponential families), the posterior update minimizes expected loss, ensuring the correction aligns with the mismatch. B1 holds by optimality, not by assumption. The gain principle $\eta^* = U_M / (U_M + U_o)$ determines only the magnitude; the direction is guaranteed by the update rule's optimality. For gradient descent on a convex loss, B1 follows from gradient monotonicity (the defining property of convexity).

Proposition B.4: Gradient Equivalence.

Statement. For any agent updating via gradient descent on a differentiable loss L with learning rate η and minimizer M^* (where $\nabla L(M^*) = 0$), there are two distinct results corresponding to the one-point and two-point sector conditions:

(B.4-i) One-point sector $\Leftarrow$ strong convexity (one direction only). The one-point sector condition at M^* — equivalent to AAT’s GA-3 / A2’ as stated in 4.2 — is

$$\delta^T F(\delta) \geq \alpha \|\delta\|^2 \quad \text{for } \|\delta\| \leq R \quad (\text{K.16})$$

evaluated at fixed equilibrium M^* . Local (α/η) -strong convexity of L on $\mathcal{B}_R(M^*)$ implies this one-point sector condition with $\alpha = \eta\mu$, but the converse fails: there exist losses satisfying the one-point sector condition that are *not* strongly convex on any neighborhood of M^* (counterexample below).

(B.4-ii) Two-point sector $\Leftrightarrow$ strong convexity (full equivalence). The two-point / incremental sector condition — DA2’-inc in O and the bridge-lemma precondition in 13.1 — is

Derived (sector-convexity-equivalence, two-point form).

$$(F(\delta_1) - F(\delta_2))^T(\delta_1 - \delta_2) \geq \alpha \|\delta_1 - \delta_2\|^2 \quad \text{for } \delta_1, \delta_2 \in \mathcal{B}_R(M^*) \quad (\text{K.17})$$

Under this strengthened condition the iff holds:

$$\text{Two-point sector with } (\alpha, R) \iff L \text{ is } (\alpha/\eta)\text{-strongly convex on } \mathcal{B}_R(M^*) \quad (\text{K.18})$$

with $\alpha = \eta\mu$ and $\mu = \inf_{\delta \in \mathcal{B}_R(M^*)} \lambda_{\min}(\nabla^2 L(M^* + \delta))$. The basin radius R is the largest ball around M^* where $\nabla^2 L$ remains positive definite.

Proof of (B.4-ii). The correction function for gradient descent is $F(\delta) = \eta \cdot \nabla L(M^* + \delta)$ (continuous-time form absorbing event rate ν into $\mathcal{T} = \nu \cdot \eta$). The two-point sector condition becomes:

$$\eta \cdot (\nabla L(M^* + \delta_1) - \nabla L(M^* + \delta_2))^T(\delta_1 - \delta_2) \geq \alpha \|\delta_1 - \delta_2\|^2 \quad (\text{K.19})$$

Dividing by $\eta > 0$ gives gradient monotonicity with modulus $\mu = \alpha/\eta$. By Nesterov 2004, Theorem 2.1.10, L is μ -strongly convex on $\mathcal{B}_R(M^*)$ iff this gradient monotonicity holds for all $x, y \in \mathcal{B}_R(M^*)$. The equivalence is bidirectional:

($\Rightarrow$) Two-point sector with (α, R) yields gradient monotonicity with $\mu = \alpha/\eta$, hence strong convexity.

($\Leftarrow$) μ -strong convexity gives gradient monotonicity, multiplying by η recovers the two-point sector condition with $\alpha = \eta\mu$. $\square$

Proof of (B.4-i) one direction. Setting $\delta_1 = \delta$, $\delta_2 = 0$ in the two-point condition and using $\nabla L(M^*) = 0$ recovers the one-point form $\delta^T F(\delta) \geq \alpha \|\delta\|^2$. So strong convexity (which is equivalent to the two-point sector by (B.4-ii)) implies the one-point sector. The reverse direction does not hold; see the counterexample below. $\square$

Counterexample showing one-point $\nRightarrow$ strong convexity. Let $L : \mathbb{R} \rightarrow \mathbb{R}$ have $L'(x) = x(1 + \frac{1}{2} \sin(10x))$ with $L'(0) = 0$ (so $x^* = 0$). The one-point sector at x^* is

$$x \cdot L'(x) = x^2 (1 + \frac{1}{2} \sin(10x)) \geq \frac{1}{2} x^2, \quad (\text{K.20})$$

so the one-point sector holds globally with $\mu = \frac{1}{2}$. But $L''(x) = 1 + \frac{1}{2} \sin(10x) + 5x \cos(10x)$ takes values such as $L''(\pi/10) = 1 - \pi/2 \approx -0.57 < 0$, so L is not convex on any neighborhood of x^* that contains $\pi/10$. The one-point sector at the equilibrium is genuinely weaker than full local strong convexity — a structural distinction, not a smoothness artifact. The full two-point sector fails on this L for any pair (δ_1, δ_2) straddling a region where $L'' < 0$ (gradient monotonicity fails there), so (B.4-ii) correctly identifies this L as a non-instance.

The α factorization (under (B.4-ii)). The decomposition $\alpha = \eta \cdot \mu$ separates agent design (η , responsiveness) from environment structure (μ , loss curvature). The adaptive tempo maps as $\alpha \propto \mathcal{T}$: higher gain or steeper curvature means faster mismatch correction.

Where each direction lands in AAT. The one-point form is what #form-sector-condition's A2' ((A2')) applied at $\delta^* = 0$ requires for Lyapunov persistence (via A) — and what AAT's six persistence-flavored results inherit through #result-sector-persistence-template (T2). The two-point / incremental form is what #form-composition-closure's bridge lemma requires (DA2'-inc) for full-update-map contraction at the composite level. Strong convexity sits at the strict end of this scale: it implies both. The one-point $\nRightarrow$ two-point counterexample above is exactly why AAT distinguishes the two — Lyapunov persistence is available for some agent classes (variational, PID-bounded-plant) where bridge-lemma contraction is not.

Loss Function Classification.

The gradient equivalence (Prop B.4) classifies which loss functions satisfy GA-3, and with what parameters:

Table K.1

Loss class	Strong convexity	α	Basin R	Status
Quadratic (Kalman, linear regression)	Global: $\mu = \lambda_{\min}(H)$	$\eta \cdot \lambda_{\min}(H)$	∞	Exact
Exponential family (natural params), bounded scope	Pointwise on the interior; uniform on any compact $\Theta_0 \subset \text{int}(\Theta)$: $\mu_0 = \inf_{\theta \in \Theta_0} \lambda_{\min}(\mathbf{I}(\theta)) > 0$	$\eta \cdot \mu_0$	Θ_0 (compact / interior-bounded)	Exact within Θ_0
L2-regularized convex loss	Global: $\mu \geq \lambda$	$\eta \cdot \lambda$ (floor)	∞	Exact
Unregularized logistic	Global convex, not strongly	$\eta \cdot \mu(R)$, decaying	∞ (but $\alpha \rightarrow 0$)	Local
Quasi-convex	$\delta^T \nabla L \geq 0$ but ratio $\rightarrow 0$	$\eta \cdot \mu(R) \rightarrow 0$	Finite effective	Local
Non-convex within basin	Local: $\mu > 0$ in basin	$\eta \cdot \mu_{\text{local}}$	Distance to nearest saddle	Local
Non-convex beyond basin	Fails: $\lambda_{\min}(\nabla^2 L) < 0$	N/A	N/A	Fails

Key observations:

1. *Exponential family models in natural parameter form satisfy GA-3 pointwise on the interior of the natural-parameter space, and uniformly on any compact subset of the interior.* The Hessian is the Fisher information matrix $\mathbf{I}(\theta)$, which is positive definite in $\text{int}(\Theta)$ — but pointwise positive definiteness does not imply a uniform global lower bound on $\lambda_{\min}(\mathbf{I}(\theta))$. The Poisson natural parameter is the canonical counterexample: $\mathbf{I}(\theta) = e^\theta$ on $\Theta = \mathbb{R}$, so $\inf_{\theta \in \mathbb{R}} \mathbf{I}(\theta) = 0$ and no global sector constant exists. A uniform $\mu_0 > 0$ — and hence a global $\alpha = \eta \mu_0$ — requires either (i) restricting the operating region to a compact $\Theta_0 \subset \text{int}(\Theta)$ on which $\mu_0 := \inf_{\theta \in \Theta_0} \lambda_{\min}(\mathbf{I}(\theta)) > 0$, or (ii) verifying a uniform Fisher lower bound for the family in question (Gaussian-mean and Beta-Bernoulli, whose Fisher information is bounded below on the interior of Θ , qualify; Poisson

and Gamma-shape do not on the full natural-parameter line). The compact-scope condition is verifiable per-application from the agent's prior support and corresponds to taking R as the diameter of Θ_0 in the natural-parameter coordinate, matching A2's local-region structure (4.2).

2. *L2 regularization transforms locally strongly convex losses into globally strongly convex ones.* The regularization parameter λ provides a floor: $\alpha \geq \eta\lambda$ everywhere.
3. *Non-convex losses satisfy GA-3 within each basin of attraction.* The basin radius R is the distance from the local minimum to the nearest inflection surface (where $\lambda_{\min}(\nabla^2 L) = 0$). Beyond the basin, the sector condition fails — this IS the structural-adaptation trigger from 4.8.

Failure Modes.

The bridge fails in precisely five cases. These characterize the exact conditions under which the gain-based update does NOT produce a correction function satisfying the sector condition.

FM-1: Directional Infidelity. The mismatch transform g rotates the correction away from the mismatch direction, so $\delta^T H g(\delta) \leq 0$.

Example. If $\delta \in \mathbb{R}^2$ and $g(\delta) = R_{90}\delta$ (a 90-degree rotation), then $\delta^T g(\delta) = 0$ for all δ . The correction is perpendicular to the mismatch and never reduces it.

When it occurs. Pathological parameterizations where the model-space update direction is misaligned with observation-space mismatch. For optimal Bayesian updates, B1 holds by construction — the posterior minimizes expected loss.

FM-2: Gain Collapse. $\eta^* \rightarrow 0$ while $\rho > 0$, so $\alpha \rightarrow 0$ and the persistence condition $\alpha > \rho/R$ eventually fails. An agent accumulating experience in a non-stationary environment will eventually have gain too small to track changes. This is not a bridge failure but a persistence-condition failure — the bridge holds ($\alpha = \eta^* \cdot c > 0$ for finite experience), but α falls below ρ/R . See 3.6.

FM-3: Nonlinear Saturation. The correction function g saturates at large $\|\delta\|$, so the sector ratio $\delta^T g(\delta)/\|\delta\|^2$ decays. For a saturating function such as $g(\delta) = \tanh(\delta)$, the local sector parameter is:

$$\alpha(R) = \eta^* \cdot \frac{\tanh(R)}{R} \quad (\text{K.21})$$

For small R , $\alpha \approx \eta^*$ (linear regime). For large R , $\alpha \approx \eta^*/R$ (saturated regime). The sector condition holds locally — the valid region shrinks as the nonlinearity strengthens. This is exactly what the local form A2' in 4.2 anticipates.

FM-4: Unobservable Directions. When $\ker(H) \neq \{0\}$ (fewer observations than state dimensions), the correction has no effect on mismatch components in $\ker(H)$. The sector condition holds only in the observable subspace $\mathcal{O} = \text{range}(H^T)$. This is a structural limitation: the agent cannot correct errors it cannot observe. See 9.3.

For persistence in the full state space, the system must be *detectable*: all unstable modes must be observable. Stable unobservable modes decay naturally without correction.

FM-5: Model Misspecification. The model class does not contain the truth, so the gradient direction is systematically wrong. B1 fails because the correction aims at the wrong target. The sector parameter degrades proportionally to the misspecification. This is the 4.8 trigger: the correction function fails not because the gain is wrong but because the correction *direction* is wrong.

Simulation Results Summary.

Six numerical experiments validate the gradient equivalence (Prop B.4) and the loss-function classification.

Experiment 1: Quadratic Loss (Global Strong Convexity). $L(w) = \frac{1}{2}(w - w^*)^T H(w - w^*)$ with $w \in \mathbb{R}^5$, Hessian eigenvalues $\{0.5, 1.0, 2.0, 3.0, 5.0\}$, $\eta = 0.1$. The sector ratio $\delta^T F / \|\delta\|^2$ at every gradient-descent step fell within $[\eta\mu, \eta L] = [0.05, 0.50]$, confirming Prop B.4 exactly. No step violated the bound.

Experiment 2: Logistic Regression (Position-Dependent Convexity). Logistic regression with $d = 2, N = 5000$ samples. Unregularized: sector ratio positive along the entire trajectory, with local μ ranging from 0.076 (at MLE) to higher values far from it. With L2 regularization ($\lambda = 0.1$): sector ratio bounded below by $\eta\lambda = 0.03$ at every point, confirming the global floor.

Experiment 3: Gaussian Mixture (Non-Convex, Basin Structure). Two-component Gaussian mixture. Within basin ($r < 3.1$): Hessian positive definite, sector ratio positive, gradient descent converges to correct optimum. At basin boundary ($r \approx 3.1$): Hessian eigenvalue crosses zero. Beyond basin ($r > 3.5$): sector ratio negative, gradient descent converges to wrong optimum. Basin boundary confirmed as sharp transition, matching the 4.8 trigger.

Experiment 4: Local Sector Parameter Measurement. Same Gaussian mixture. Sector parameter $\alpha(r)$ decreases monotonically from 181.4 (at $r = 0.5$) to 125.8 (at $r = 3.0$), then drops to zero at $r \approx 3.1$. The decay is moderate within the basin (69% retention at the boundary); the transition at the boundary is sharp.

Experiment 5: Exponential Family (Poisson, Natural Parameters). Poisson in natural parameter space. Sector ratio positive at every step, ranging from 0.00177 to 0.00240 (with $\eta = 0.001$). Confirms global strong convexity of exponential family losses in natural parameterization.

Experiment 6: Quasi-Convex and Multi-Modal Losses. *Quasi-convex* ($L(w) = -1/(1 + w^2)$): sector ratio always positive but decays as $O(1/w^4)$. No uniform lower bound exists — $\alpha(R) \rightarrow 0$, making persistence bounds vacuous for large R .

Multi-modal ($L(w) = \sin(w) + 0.01w^2$): sector condition holds within each basin ($|\delta| < 3.0$), fails at basin boundary (gradient reversal at local maximum). Basin radius $\approx$ distance from local minimum to nearest local maximum.

Epistemic Status.

Propositions B.1 and B.2 are *exact* — they follow from standard Kalman filter algebra. The only subtlety is the weighted norm in the matrix case (B.2), which is a quantitative refinement, not a qualitative gap. Both are standard results in estimation theory reframed in sector-condition language.

Proposition B.3 (Bridge Theorem) is a *conditional derivation*: exact under B1 (directional fidelity). B1 holds by construction for optimal Bayesian updates (the posterior minimizes expected loss, ensuring the correction aligns with the mismatch). For approximate update rules, B1 is a design condition, not a global assumption. The condition is transparent and checkable for specific systems.

Proposition B.4 (Gradient Equivalence) splits along the one-point / two-point distinction. (B.4-ii) — the iff between strong convexity and the *two-point / incremental* sector condition (DA2'-inc) — is an *exact mathematical equivalence*: bidirectional via Nesterov 2.1.10. (B.4-i) — the relation to AAT's one-point sector A2' as stated in 4.2 — is *one-directional only*: strong convexity implies the one-point sector at M^* , but not conversely. The Codex-style counterexample $L'(x) = x(1 + \frac{1}{2}\sin(10x))$ satisfies the one-point sector globally yet has L'' negative on intervals, ruling out a converse. The loss-function classification table is read accordingly: rows that supply strong convexity satisfy both forms; rows that supply only one-point sector (none in the current table — strong convexity is the universal column)

would not transfer to two-point. The simulation experiments confirm the theoretical predictions across all six loss classes tested; no violations were observed where the theory predicts the sector condition holds.

Max attainable: *conditional* for the bridge (B1 is inherent — pathological update rules exist), *exact* for the gradient equivalence in its two-point form (B.4-ii) and *exact one-direction* for its one-point form (B.4-i). The condition cannot be removed: there exist correction functions that violate the sector condition (FM-1 through FM-5), and there exist losses that satisfy the one-point sector at M^* without being strongly convex (the Codex-style counterexample), which is why the one-point $\nRightarrow$ two-point gap is structural rather than a smoothness artifact.

What remains open. (1) Non-gradient agents (PID controllers, rule-based systems, human judgment): GA-3 remains an empirical claim for these agent classes. (2) Time-varying α : when the gain adapts (Adam, RMSprop, Kalman transient), $\alpha(t)$ varies; the Lyapunov analysis extends to time-varying $\alpha(t) \geq \underline{\alpha} > 0$ via standard results (Khalil 2002, Chapter 8). (3) Stochastic gradients: SGD satisfies the sector condition *in expectation*; the per-step noise enters as effective disturbance in the Prop A.1S framework.

Discussion.

The formal chain is closed for well-designed agents. The prediction chain becomes:

gain principle+B1 (directional fidelity) $\xrightarrow{\text{Prop B.3 (exact)}}$ sector condition (GA-3) $\xrightarrow{\text{Props A.1, A.1S, A.2 (exact)}}$ persistence, reserve, scaling (K.22)

For optimal Bayesian updates, B1 holds by construction, making the full chain a derivation. For gradient agents, B1 is equivalent to local strong convexity — a well-characterized property.

GA-3 is grounded, not floating. The assumption load shifts from “the correction function satisfies the sector condition” (what correction function?) to “the update rule has directional fidelity” (the correction at least points toward reality). For the important agent classes — Kalman filters, conjugate posteriors, exponential families, gradient descent on convex losses — this is automatic.

Basin boundary = structural adaptation trigger. For gradient agents with non-convex losses, the basin radius R is the convexity radius of the loss landscape. When mismatch exceeds R , the correction function reverses direction. This is the 4.8 trigger with a precise geometric characterization: structural adaptation is needed when the agent crosses an inflection surface of its loss landscape.

Working Notes.

- The fluid-limit gap: the bridge analysis works in expected value. The actual per-step correction is stochastic; observation noise ε contributes $K\varepsilon$ to each update, acting as effective disturbance of magnitude $\|K\|\sigma_\varepsilon$ per step. This maps to the Prop A.1S (stochastic) framework — the sector condition on the expected correction plus stochastic noise.
- For variational inference and other approximate methods, B1 is not guaranteed by optimality and must be verified per approximation scheme.
- The adaptive-reserve factorization for gradient agents is $\Delta\rho^* = \eta\mu R - \rho$, with three controllable factors: gain η , curvature μ , and basin width R . An agent is robust when $\eta\mu R \gg \rho$.
- Full simulation code, raw traces, and intermediate analysis live in `spikes/.integrated/spike-gain-sector-bridge-nonlinear.md`; the segment carries the parameters and outcomes that establish the result. Landing-context provenance: this segment consolidates `spikes/.integrated/spike-gain-sector-bridge.md` (Kalman / bridge theorem) and `spikes/.integrated/spike-gain-sector-bridge-nonlinear.md` (gradient equivalence and the six experiments).

Incidental audit gold (gold-lift, 2026-05-31). Cross-audit ideation harvested from de-novo auditors’ working dirs, deduplicated and lightly attributed. Orthogonal framing / pedagogy material; off-ramp (already-resolved + dimensional-convention) at the end. **Coverage:** three substrates reached a digested reflection (Gemini, AUDIT-WORKING-193847; Claude, AUDIT-WORKING-584721; Codex, AUDIT-WORKING-742613), plus cross-cycle confirmation from Claude, AUDIT-WORKING-963715 and Gemini, AUDIT-WORKING-849201.

1. *Candidate Briefprose / pre-prose.*

- **“The Rosetta Stone between Control and ML.”** Prop B.4’s gradient-equivalence (under the corrected one-way reading) shows gradient descent on a strongly-convex loss is the same object as a Kalman filter on a linear system — both instances of the sector condition. A plain-language anchor for why the bridge matters: it lets control-theoretic persistence bounds apply to learning systems (Gemini, AUDIT-WORKING-193847).
- **The $\alpha = \eta\mu$ factorization as the diagnostic.** It separates *agent-design* (gain η , tunable by the designer) from *environment-structure* (curvature μ , basin width R , given by the world); adaptive reserve $\Delta\rho^* = \eta\mu R - \rho$ then reads as a three-factor budget, and “robustness = $\eta\mu R \gg \rho$ ” is the one-line takeaway (Claude, AUDIT-WORKING-584721).

2. *Candidate Discussion.*

- **The FM-1–FM-5 failure-mode taxonomy as a diagnostic for learning collapse.** Two substrates singled out the failure-mode classification as standout pedagogy — it formally separates vanishing-gradient (FM-3), misspecification (FM-5), and unobservable-latent-state (FM-4) failures. Candidate to surface more prominently as the segment’s diagnostic contribution (Gemini, AUDIT-WORKING-193847; Claude, AUDIT-WORKING-584721 praises the loss-classification table as a model for other segments).
- **Directional infidelity (FM-1) as organizational “blame-shifting.”** When the update rule rotates the correction away from the true mismatch direction, $\delta^\top g(\delta) \rightarrow 0$ and the sector parameter $\alpha \rightarrow 0$ even at high gain η^* : the system expends huge effort and learns nothing. The organizational analog (firing a mid-level manager instead of fixing the structural flaw) makes the abstract orthogonality condition concrete, and the consciousness-infrastructure reading — *the internal update must point at the source of the error; an agent that rotates its own error-gradient to protect a favored hypothesis suffers FM-1* — is a vivid pedagogical extension (Gemini, AUDIT-WORKING-193847).
- **Quasi-convex plateaus are worse than steep narrow basins.** The $\alpha \rightarrow 0$ -as- $R \rightarrow \infty$ case: a plateau gently sloping toward a hole provides *zero adaptive reserve* against drift, whereas a narrow steep non-convex basin at least gives strong local correction — a counterintuitive consequence of the sector geometry worth naming (Gemini, AUDIT-WORKING-193847).

3. *Follow-up items.*

- **The “expected-bridge + stochastic-disturbance” composition is a meta-pattern.** The bridge analysis works in expected value (already in Working Notes as the fluid-limit gap); the stochastic per-step residual $\|K\|\sigma_\varepsilon$ enters as Model-S disturbance under Prop A.1S — so the framework’s expected-value bridge results compose cleanly with its stochastic noise treatments. Flagged as a candidate cross-segment structural observation, not just a local note (Claude, AUDIT-WORKING-584721).
- The fully-observable diagonal case $\alpha = \min_i K_i$ (bottleneck dimension sets persistence) lines up with #result-per-dimension-persistence; the 72% scalar-overestimate from #def-adaptive-tempo anisotropy presumably comes from averaging where minimization is correct (Claude, AUDIT-WORKING-584721).

4. *Readers often ask / wonder.*

- **“Is the gradient equivalence really an iff?”** This is exactly the question the corrected B.4 (one-way $\Leftarrow$) answers — see off-ramp; worth a one-line forward-pointer so a fresh reader does not re-derive the false converse (Codex, AUDIT-WORKING-742613; verified first-hand at Claude, AUDIT-WORKING-963715 — “mathematically precise and consequential”).

Off-ramp (NOT gold) — routed for adjudication.

- **Prop B.4 gradient-equivalence “iff” overstatement — RESOLVED-BY-STRENGTHENING** (recorded so it is **not re-opened**). Codex (AUDIT-WORKING-742613, “candidate finding K”) and Gemini (AUDIT-WORKING-849201) independently flagged that the one-point sector condition (monotonicity relative to M^*) is *strictly weaker* than local strong convexity; the proof’s reverse direction silently replaced the one-point inequality with full two-point gradient monotonicity (witness $L'(x) = x(1 + \frac{1}{2} \sin(10x))$: one-point sector holds with $\mu = \frac{1}{2}$ while L'' goes negative nearby). Per the cycle records (742613-F3, resolved by strengthening; landing at the B.4 split — one-point $\Leftarrow$ vs two-point $\Leftrightarrow$), this was **already corrected in current canon**. Flagged here only as a regression-guard: the segment’s counterexample $L'(x) = x(1 + \frac{1}{2} \sin(10x))$ is the *defense* of the corrected one-way reading, not a defect — do not “restore” an iff.
 - **Exponential-family global-strong-convexity scope (likely subsumed by the B.4 correction)**. Codex (AUDIT-WORKING-742613): “all exponential-family models in natural-parameter form satisfy GA-3 globally” is too strong — pointwise PD Fisher information does not give a uniform global lower bound (Poisson natural parameter: Fisher = e^θ , $\inf_{\mathbb{R}} = 0$). Global sector constants need a compact / bounded-away-from-boundary region or a uniform Fisher lower bound. Verify whether the B.4 split already scoped this.
 - **Per-event vs per-time normalization (architectural; tracked in PROPOSALS)**. The deeply-mathematical AUDIT-WORKING-526815 cycle (F13) flags that #der-gain-sector-bridge derives $\alpha = \eta^* c_{\min}$ while persistence uses α as a *rate* and tempo is $\mathcal{T} = \nu \eta^*$ — the event-rate factor ν is omitted unless F is time-aggregated. Cross-cycle convergence (471203, 526815) reads this as “the bridge is *incomplete*, and downstream segments treat it as complete,” with the strengthen-first answer being to finish the bridge + its proof dependencies (this segment, #deriv-sector-condition) *first*, then propagate the cleaned dimensional convention. Routed architectural $\rightarrow$ PROPOSALS, connected to PRACTICA’s persistence-and-stability area.
 - **FM-3 saturation hard ceiling (scope-sharpening)**. Gemini (AUDIT-WORKING-193847): for $g(\delta) = \tanh(\delta)$, the local sector parameter $\alpha(R) = \eta^* \tanh(R)/R$ correctly decays, but the *force* saturates — so survival is impossible if $\rho > \|F\|_{\max}$, independent of basin radius R or linear-regime gain η^* . Suggested: add a note to FM-3 that saturating nonlinearities cap the maximum absolute correction speed.
-

Appendix L

Derivation: Recursive Update — Uniqueness Derivation

Appendix derivation backing 3.2. The claim being established is *uniqueness*: under three constraints — arrow of time, partial observability, and state completeness — the recursive form $M_{\tau+} = f(M_{\tau-}, e_{\tau})$ is the only update form realizable, not merely one consistent option. The structure is an elimination argument over the universe of information potentially accessible at event time τ : the environment state, the full chronica, all prior model states, the current event e_{τ} , and future events.

The three constraints do different kinds of work. **C1 (arrow of time, from 1.7)** rules out dependence on future events: in a classical universe such information is simply not accessible, and predictions about the future, however accurate, are internal computations belonging to $M_{\tau-}$, not future information. **C2 (partial observability, from 1.5 and 1.3)** rules out direct access to Ω_{τ} except through the lossy observation channel e_{τ} . **C3 (state completeness, from 2.1)** compresses all prior-past information — the full chronica and every earlier model state — into $M_{\tau-}$, because by construction M is the agent’s complete internal state and any information available to the update but not in M would mean M was misspecified. After all three eliminations only $(M_{\tau-}, e_{\tau})$ remains in the dependency set. A measure-theoretic version via the Doob–Dynkin lemma — restricting the agent’s sigma-algebra to $\sigma(M_{\tau-}, e_{\tau})$ — gives a cleaner technical proof.

The derivation then mounts seven adversarial attacks to test whether the result really holds. Simultaneous events are absorbed by treating the bundle as a single e_{τ} (form preserves). Continuous environmental coupling is genuinely outside the event-driven formulation — but the analogous result reached by the same three-constraint argument in continuous time is the state-space representation $\dot{M} = g(M, u)$ from classical control theory, so the underlying structure survives. The objection that C3 is circular — anything not in M gets absorbed into M by definition — is acknowledged as the deepest objection, and the derivation is explicit about it: the real content is an *analytical commitment* that by defining M as complete one commits to Markovian analysis, which then makes 2.3 the right quality metric. Shared state, external randomness, time-dependence, and agents that store full history each reduce to the recursive form by appropriate enlargement of the model space or interpretation of f .

The honest epistemic accounting at the end is load-bearing. **C1** (physical postulate) and **C2** (scope definition) do *eliminative* work — they rule out physically impossible and scope-excluded update forms. **C3** (analytical commitment) does not eliminate; it is *definitional* and cannot be “violated” because any apparent violation is absorbed by expanding M . The recursive form’s uniqueness is therefore conditional on accepting the three-constraint set, not on the constraints being independently inescapable — C3 in particular *could* be refused, at the cost of leaving AAT’s scope and accepting non-Markovian analysis. The Markov property is not a discovered law of adaptive systems but the unique form compatible with the modeling commitment of 2.1.

Setup.

We work within AAT's scope (1.5): an agent coupled to an environment Ω through observation and action channels, with residual uncertainty.

Universe of information at event time τ . The following information exists (in the broadest ontological sense) at the moment event e_τ occurs:

Table L.1

<i>Information</i>	<i>Description</i>
Ω_τ	The environment state
$\mathcal{C}_{\tau-}$	The complete interaction history (1.4) up to (but not including) e_τ
$\{M_{\tau'}\}_{\tau' \leq \tau-}$	The agent's prior internal states, culminating in $M_{\tau-}$
e_τ	The current event (observation arriving or action completing)
$\{e_{\tau'}\}_{\tau' > \tau}$	Future events (not yet occurred)

The question: of these, which can the update $M_{\tau+}$ depend on?

The Three Constraints.

Constraint 1 — Arrow of time (1.7 postulate). Events are temporally ordered and this ordering is irreversible. An update occurring at time τ cannot depend on events that have not yet occurred:

$$M_{\tau+} \text{ cannot depend on } \{e_{\tau'}\}_{\tau' > \tau} \quad (\text{L.1})$$

This is a physical constraint — the most primitive one. In a classical universe, information from the future is simply not available. Even if the agent can *predict* future events, those predictions are part of $M_{\tau-}$ (they are internal computations, not future information).

Constraint 2 — Partial observability (1.5). The agent cannot access Ω_τ directly. Its only interface with the environment is through the event e_τ , which is a lossy function of Ω_τ (via 1.3):

$$M_{\tau+} \text{ cannot depend on } \Omega_\tau \text{ except through } e_\tau \quad (\text{L.2})$$

This is a scope constraint. If the agent could access Ω directly, the residual uncertainty condition in 1.5 would be trivially violable.

Constraint 3 — State completeness (2.1). $M_{\tau-}$ is the agent's *complete* internal state just before event e_τ . There is no information about the agent's past that is available to the update mechanism but not encoded in $M_{\tau-}$:

$$M_{\tau+} \text{ cannot depend on } \mathcal{C}_{\tau-} \text{ or } \{M_{\tau'}\}_{\tau' < \tau-} \text{ except through } M_{\tau-} \quad (\text{L.3})$$

This constraint does the most interesting work and deserves careful examination (see Discussion below).

The Derivation.

Result (Recursive Update Uniqueness). Under Constraints 1–3, the model update at event time τ must have the form

$$M_{\tau+} = f(M_{\tau-}, e_{\tau}) \quad (\text{L.4})$$

for some (possibly stochastic) function $f : \mathcal{M} \times \mathcal{E} \rightarrow \mathcal{M}$. No other update form is consistent with the three constraints.

Derivation. Consider the most general possible update. The updated state $M_{\tau+}$ is a function of *all accessible information*:

$$M_{\tau+} = F(\text{accessible information at } \tau) \quad (\text{L.5})$$

We characterize the accessible information by eliminating what is not accessible.

(i) Eliminate future events. By C1 (arrow of time), $\{e_{\tau'}\}_{\tau' > \tau}$ is not accessible.

After this elimination, the candidate dependency set is: $\{\Omega_{\tau}, \mathcal{C}_{\tau-}, \{M_{\tau'}\}_{\tau' \leq \tau-}, e_{\tau}\}$

(ii) Eliminate direct environment access. By C2 (partial observability), the agent cannot access Ω_{τ} except through the event e_{τ} . Any information from Ω_{τ} that reaches the agent does so through e_{τ} — already in the dependency set.

After this elimination: $\{\mathcal{C}_{\tau-}, \{M_{\tau'}\}_{\tau' \leq \tau-}, e_{\tau}\}$

(iii) Reduce past information to $M_{\tau-}$. By C3 (state completeness), $M_{\tau-}$ is the agent's complete internal state. Every element of $\mathcal{C}_{\tau-}$ and every prior model state $M_{\tau'}$ ($\tau' < \tau-$) that could influence the update can do so *only through* its effect on $M_{\tau-}$. The agent's internal state evolves through a sequence of updates; the cumulative effect of all prior events is exactly $M_{\tau-}$. The raw events that produced this state are no longer separately available — they were “consumed” by the update mechanism and their information (to the extent it was retained) is now encoded in $M_{\tau-}$.

Could the agent maintain a separate log of raw events outside of M ? It could — but that log *is part of* M . Whatever information the agent retains in any form — model parameters, cached data, raw event buffers, metadata — is by definition part of its complete internal state $M_{\tau-}$. If something is available to the update mechanism and not in $M_{\tau-}$, then $M_{\tau-}$ was not the complete state — contradicting C3.

After this elimination: $\{M_{\tau-}, e_{\tau}\}$

Therefore: $M_{\tau+} = F(M_{\tau-}, e_{\tau}) \equiv f(M_{\tau-}, e_{\tau})$

This is the unique form: no information beyond $(M_{\tau-}, e_{\tau})$ is accessible under the three constraints, so no update form depending on anything else is realizable. $\square$

Corollary (Between-events dynamics). Between events, no new event e arrives. The same argument applies with e_{τ} removed from the accessible set:

$$\frac{dM}{d\tau} = g(M_{\tau}) \quad (\text{L.6})$$

The agent's internal evolution between events (prediction, decay, internal simulation) depends only on the current state. $\square$

Corollary (Serial special case). When observations and actions alternate at a uniform rate on a single channel, each event e_t is the pair (o_t, a_{t-1}) . The update becomes:

$$M_t = f(M_{t-1}, o_t, a_{t-1}) \quad (\text{L.7})$$

This is the familiar discrete-time form. $\square$

Information-Set Formalization.

For readers who prefer a measure-theoretic framing:

The agent's **information set** at time τ is the sigma-algebra $\mathcal{J}_\tau^{\text{agent}}$ — the collection of events (in the probability-theoretic sense) about which the agent can condition its update.

- **C1** restricts $\mathcal{J}_\tau^{\text{agent}} \subseteq \sigma(\{e_{\tau'} : \tau' \leq \tau\} \cup \{\Omega_\tau\} \cup \{M_{\tau'} : \tau' \leq \tau^-\})$ — no future information.
- **C2** further restricts: $\sigma(\Omega_\tau) \setminus \sigma(e_\tau)$ is not in $\mathcal{J}_\tau^{\text{agent}}$ — the agent cannot condition on aspects of Ω_τ not captured by e_τ .
- **C3** further restricts: $\sigma(\{e_{\tau'} : \tau' < \tau\} \cup \{M_{\tau'} : \tau' < \tau^-\}) \subseteq \sigma(M_{\tau-})$ from the agent's perspective.

After all three restrictions: $\mathcal{J}_\tau^{\text{agent}} = \sigma(M_{\tau-}, e_\tau)$.

By the Doob–Dynkin lemma¹, any $\sigma(M_{\tau-}, e_\tau)$ -measurable random variable is a (Borel) function of $(M_{\tau-}, e_\tau)$. Therefore $M_{\tau+} = f(M_{\tau-}, e_\tau)$ for some measurable f . $\square$

Attempts to Break the Result.

Before trusting the proof, seven counterexample attacks:

Attack 1: Simultaneous events. Two events arrive at exactly the same time: $e_\tau^{(1)}$ and $e_\tau^{(2)}$. The update has three arguments: $f(M_{\tau-}, e_\tau^{(1)}, e_\tau^{(2)})$.

Verdict: Not deep — 3.1 defines events as atomic. If we allow bundled events, the form holds with e_τ as a set. Reveals that “event” needs careful definition, but the form is preserved.

Attack 2: Continuous environmental influence. An agent embedded in a physical system experiences continuous forces (gravity, temperature, electromagnetic fields). These aren't “events” in 3.1's sense; they're continuous signals. The true dynamics would be $dM/d\tau = g(M_\tau, o(\tau))$ where $o(\tau)$ is a continuous observation stream.

Verdict: Genuine limitation of the event-driven formulation. The between-events corollary $dM/d\tau = g(M_\tau)$ holds only when the agent is truly isolated between events. For continuous coupling, the analogous result is the general state-space representation $\dot{M} = g(M, u)$ from control theory — arrived at by the same three constraints. The event-driven version is a special case for digital/sampled systems.

Attack 3: The C3 circularity. C3 defines M as the agent's complete internal state. Any apparent counterexample is dissolved by expanding M . Consider: an agent has a “model” (neural net weights) and a “replay buffer” (stored raw events). C3 says $M = (\text{weights}, \text{buffer})$. The model space is just larger than you thought.

Verdict: The deepest objection. The proof essentially: (1) Define M to be everything the agent has. (2) Observe the update can only use what the agent has. (3) Therefore $f(M_{\tau-}, e_\tau)$. The real content is the *analytical commitment*: by defining M as complete, we commit to Markovian analysis, which then makes 2.3 the right quality metric. See Epistemic Status below.

¹Kallenberg, O. (2002). *Foundations of Modern Probability* (2nd ed.). Springer. §1.2 (measurability and the Doob–Dynkin lemma).

Attack 4: Shared state between agents. Agents A and B share a common memory bank (shared database). The clean resolution is the multi-agent framework: the shared memory is part of the *composite* system’s state, and each agent’s interaction with it is mediated by events (reads and writes). Not a true counterexample but highlights that C3 requires careful delineation of agent boundaries.

Attack 5: External randomness not in e_τ . Hardware thermal noise used in the update. The stochastic case $M_{\tau+} \sim P(\cdot \mid M_{\tau-}, e_\tau)$ is a special case of f where f is a randomized function. The *form* — dependence on exactly $(M_{\tau-}, e_\tau)$ — is preserved. The result statement should explicitly allow stochastic f .

Attack 6: Time-dependent updates. Could f depend on the timestamp τ itself? Yes — consistently. The event e_τ in 3.1 carries a timestamp: $e_\tau = (\text{type}, \text{channel}, \text{payload}, \tau)$. So time-dependence enters through e_τ . Alternatively, the agent may maintain an internal clock as part of $M_{\tau-}$. Either way, $f(M_{\tau-}, e_\tau)$ accommodates time-dependence.

Attack 7: Agents that store full history. An agent with $M_{\tau-} \supseteq \mathcal{C}_{\tau-}$ is entirely consistent. The model space $\mathcal{M}$ is simply large enough to include the raw history. The 2.2 argues compression is *wise* — but the recursive update form holds regardless of compression level.

What Is Derived vs. What Is Chosen.

Table L.2

<i>Property</i>	<i>Source</i>	<i>Strength</i>
Constraint C1 (arrow of time: update depends on τ^- , not future events)	Physical law — not a formulation choice	Postulate (physical)
Constraint C2 (partial observability: update depends on e_τ , not raw Ω_τ)	Scope definition of AAT	Postulate (scope-defining)
Constraint C3 (state completeness: $M_{\tau-}$ summarizes the agent’s relevant past)	Analytical commitment — the definition of M as complete	Definition
Recursive form $M_\tau = f(M_{\tau-}, e_\tau)$	C1 + C2 + C3	Proved (unique form compatible with the three constraints)
Future-dependent updates eliminated	C1 alone	Derived (direct consequence)
Ω_τ -dependent updates eliminated	C2 alone	Derived (direct consequence)
Full-history-dependent updates reducible to recursive form	C3 + any choice of $M \supseteq \mathcal{C}_{\tau-}$	Proved (compatibility, not elimination)
Markov property of the update	C3 (completeness) + recursive form	Proved (follows from C3 definition)
Seven attack counterexamples (simultaneous events, continuous coupling, C3 circularity, shared state, external randomness, time-dependence, full history)	Case-by-case reduction to the recursive form	Proved (each)
C3 is definitional, not eliminative	Analysis of what C3 asserts vs. what it rules out	Discussion-grade (clarifying observation)

The dividing line: C1 and C2 do genuine *eliminative* work — they rule out physically or scope-excluded update forms. C3 is a *definitional commitment* that forces the Markov structure by making M complete by construction; it cannot be “violated” because any apparent violation means M was misspecified. The recursive form’s uniqueness is therefore conditional on the three-constraint set being accepted, not on the constraints being independently inescapable — C3 in particular could be refused (yielding non-Markovian analysis), at the cost of leaving AAT’s scope.

Epistemic Status.

The result is correct but partly definitional. The three constraints have different epistemic characters:

Table L.3

<i>Constraint</i>	<i>Character</i>	<i>Can it be violated?</i>
C1 (arrow of time)	Physical law	Not in a classical universe
C2 (partial observability)	Scope definition	Only by leaving AAT's scope
C3 (state completeness)	Analytical commitment	Not without redefining M

C1 and C2 do genuine eliminative work — they rule out update forms that depend on future events or on raw Ω . These are non-trivial constraints.

C3 is a definitional commitment that produces the Markov structure. It cannot be “violated” because any violation is absorbed by expanding M . This is not a weakness — it’s the nature of the claim. The result says: *the Markovian analysis is the only one consistent with C1 + C2 + the definition of M as complete*. The alternative — an update that depends on something outside M — is not “wrong” but rather means M was misspecified.

What the result says: C1 eliminates a physically impossible class of updates (future-dependent). C2 eliminates a scope-excluded class (Ω -dependent). After (1) and (2), the *only remaining question* is how the past enters: through the full history $\mathcal{C}_{\tau-}$ or through a compressed state $M_{\tau-}$. C3 says the agent *has* a complete state, and whatever that state is, it’s all the agent has. The Markov form follows.

What the result does NOT say: That M must be a lossy compression (the agent could store full history). That the Markov property is “natural” or “optimal” (it’s a consequence of how M is defined). That continuous-coupling systems are event-driven (the event framework is one abstraction; $\dot{M} = g(M, u)$ is the more general one, arrived at by the same three constraints).

Discussion.

Recursion as a consequence of completeness. The recursive form is not an assumption bolted on — it follows from the definition of M_t as complete. The sufficiency of the recursive form is precisely what 2.3 measures: when $S(M_t) = 1$, the recursive update loses nothing.

What this opens. The proof yields the *form*. It immediately invites the follow-up questions that the rest of the theory addresses: What should f preserve? → 2.2 and 2.3. How should f weight new information? → 3.6. When is $\mathcal{M}$ itself inadequate? → 4.8.

Working Notes.

- C3’s definitional character is a feature, not a bug — but it must be stated honestly. The result is not “the update must be Markovian” but rather “the Markovian analysis is the *only* consistent one, given the modeling commitment of 2.1.” These sound the same but have different epistemic status.
- The continuous-coupling generalization (Attack 2) deserves a proper note somewhere: $\dot{M} = g(M, u)$ is the more general form, with event-driven updates as a special case. The three constraints produce the same argument structure in both cases.
- The information-set formalization (Doob-Dynkin) provides the cleanest technical proof. It should probably be considered the primary proof path, with the elimination argument as the more intuitive exposition.

Incidental audit gold (gold-lift, 2026-05-31). Cross-audit ideation harvested from de-novo auditors' working dirs, deduplicated and lightly attributed. Orthogonal framing / pedagogy; off-ramp (status-label harmonization) at the end. **Coverage:** five substrates reached a digested reflection (Gemini, AUDIT-WORKING-193847; Claude, AUDIT-WORKING-584721; Claude, AUDIT-WORKING-471203; Codex, AUDIT-WORKING-742613; Gemini-voiced, AUDIT-WORKING-773921) — unusually strong convergence; several rated this the most-trusted derivation they had read.

1. Candidate Brief prose / pre-prose.

- **“The Markov property is a boundary you draw around the agent, not a fact about the world.”** The recurring one-line synthesis across substrates: most ML/RL frameworks *postulate* the Markov property; here it is *forced* by C1 (arrow of time) + C2 (partial observability) + C3 (state completeness), with C3 the definitional commitment. “A rigorous proof that the Markov assumption in RL is not an assumption about the world, but a tautological boundary drawn around the agent” (Gemini-voiced, AUDIT-WORKING-773921; same framing, Gemini, AUDIT-WORKING-193847; Claude, AUDIT-WORKING-471203).
- **“Memory as digestion.”** When event e_τ arrives, the agent extracts the update and discards the raw event; M_t is “the accumulated nutritional value of all past events,” not a storage bin. A concrete everyday analog for the recursive form that a non-specialist can re-derive (Gemini, AUDIT-WORKING-193847).

2. Candidate Discussion.

- **This derivation as the template for AAT’s inevitability-core segments.** Three substrates independently read the segment’s structure — three named constraints (eliminative C1/C2 vs definitional C3), dual proof paths (set-elimination + Doob-Dynkin), seven explicit counterexample “attacks” with honest verdicts, and the “What Is Derived vs. What Is Chosen” table — as the form *every other* AAT inevitability claim should adopt (Claude, AUDIT-WORKING-584721 and AUDIT-WORKING-471203; the seven-attack discipline named as a candidate meta-pattern, e.g. #result-mismatch-decomposition could enumerate “what would break the bias-variance decomposition?”).
- **“Epistemic-architectural rather than mathematical.”** Naming-seed for AAT’s distinctive contribution as legible here: most frameworks contribute new math; AAT’s distinctive move may be new *forms of stating* what is known, with explicit constraint-naming, definitional-character disclosure, and attack-defense. Closer to a philosophy-of-science contribution than a purely mathematical one — candidate for framing-level material (Claude, AUDIT-WORKING-471203). (*Note the early finding-vs-framing texture: offered as a re-characterization of the whole framework’s contribution, not just this segment.*)

3. Follow-up items.

- **The C3-circularity acknowledgment must survive downstream.** Several auditors flag a consistency-watch: segments that *use* the recursive-update result should preserve the “C3 is definitional, the result is conditional-on-the-modeling-commitment” honesty rather than treating Markovianity as physically derived; drift here is a cross-segment finding (Claude, AUDIT-WORKING-471203; Codex, AUDIT-WORKING-742613).
- **Continuous-coupling caveat propagation.** Attack 2’s $\dot{M} = g(M, u)$ generalization is the more-general form; the main #der-recursive-update between-event corollary $dM/d\tau = g_M(M_\tau)$ should probably carry the same “event-driven is the special case” caveat more prominently (Codex, AUDIT-WORKING-742613; already noted in this segment’s own Working Notes above — the auditors converged on it independently).
- **The “(Descended from TFT Appendix ...)” annotation is diff-voice in the body.** Editorial: move to Working Notes or remove; flagged as the fourth instance of the diff-voice pattern across Section I (Claude, AUDIT-WORKING-471203).

4. Readers often ask / wonder.

- **“What is the discretization penalty for the continuous case (Attack 2)?”** If the true dynamic is $\dot{M} = g(M, o(\tau))$, the event-driven form is a Riemann-sum-style approximation — does the agent suffer an integration error scaling with the inter-event gap? (Gemini, AUDIT-WORKING-193847).

- “Doesn’t a lossy ϕ make the recursive form epistemically sub-optimal?” If $M_{\tau-}$ is not a sufficient statistic for $\mathcal{C}_{\tau-}$ (lossy compression per #form-information-bottleneck), the recursive form locks the agent out of re-querying discarded history — mathematically valid but a structural disadvantage versus an agent that can read raw $\mathcal{C}_{\tau-}$. Offered as the formal motivation for *why* agents build external memory systems (writing things down bypasses the recursive bottleneck) (Gemini, AUDIT-WORKING-193847). (*Connects to SP-27 / lossy- ϕ recursion and #disc-m-preservation.*)

Belongs elsewhere.

- **Logogenic / Section IV reading** (points at 03-llm-core/ #disc-m-preservation, not this appendix): the C3-as-definitional move means a logogenic agent with externalized M_t (on disk) can have its M_t *defined* to include the externalized state, and the recursive form holds — the formal substrate for the M_t -preservation argument. Also the provocative architectural claim that Transformers “carry the raw apples around” (re-processing the full context each forward pass) rather than digesting history, so an external recursive loop must extract M_t and feed it forward — aspirational reach toward the logogenic-wrapping material (Gemini, AUDIT-WORKING-193847; Claude, AUDIT-WORKING-471203 names #disc-m-preservation as the destination).

Off-ramp (NOT gold) — routed for adjudication.

- **Status-label harmonization: exact (this appendix) vs conditional (#der-recursive-update body) — a genuine cross-substrate disagreement worth preserving as signal.** Two readings landed: (a) Claude (AUDIT-WORKING-471203) initially flagged a mismatch, then **withdrew** it after reading — concluding the layering is *deliberate and honest*: the appendix is “exact given the three constraints,” the body is “conditional on accepting C3 as a modeling commitment,” two angles on the same derivation. (b) Codex (AUDIT-WORKING-742613) and Gemini-voiced (AUDIT-WORKING-773921) judged the *language* should still be harmonized, proposing “exact conditional on C1–C3 and the event-driven/sampled representation; continuous coupling generalizes the input form.” Net: not a defect in the math, possibly a wording-clarity improvement on whether the two statuses read as obviously-the-same-claim-at-two-layers. Light-touch; route as a status/voice-clarity check, not a correctness finding.

Appendix M

Derived: Multi-Timescale Stability

When adaptive processes operate at nested timescales, the composite is stable exactly when each level satisfies the sector-persistence template with two interconnection terms added to its effective disturbance — the residue of the settled level below, and the target-drag of the moving level above — and the qualitative convergence constraint $\nu_{n+1} \ll \nu_n$ (4.9) becomes a closed-form threshold: the admissible timescale ratio equals the faster level's adaptive reserve divided by the rate at which the slower level drags the faster level's target.

A mental model before the formalism: picture each adaptive level as a bathtub — water is that level's mismatch, the drain is its correction, and the inflow is its disturbance (4.5 carries the single-tub version). Stacking levels couples the tubs in two directions. When the slow level acts, it *moves the fast level's drain* — the fast tub must chase a shifting target, which behaves exactly like extra inflow. And the fast tub's residual sloshing *splashes into the slow tub* — the slow level's measurements are contaminated by however much the fast level has not yet settled. Both couplings are ordinary disturbance once named, so each tub needs only the ordinary persistence condition against its *total* inflow — and the theorem's two conditions say precisely that. Micromanagement is the first coupling overwhelming the fast tub (the target moves faster than the fast level's spare drain capacity); catastrophic forgetting is the second coupling overwhelming the slow tub (the slow level is too sensitive to fast-level splash). They are dual violations of one theorem.

Formal Expression.

Fast state $x_1 \in \mathbb{R}^{n_1}$, slow state $x_2 \in \mathbb{R}^{n_2}$:

Formulation (stacked two-timescale system).

$$\dot{x}_1 = -F_1(x_1; x_2) + w_1(t), \quad \dot{x}_2 = \epsilon [-F_2(x_2; x_1) + w_2(t)] \quad (\text{M.1})$$

with $\epsilon > 0$ the timescale ratio — in AAT vocabulary $\epsilon = \nu_2/\nu_1$, the adjacent-level event-rate ratio whose smallness 4.9 asserts qualitatively. The singular-perturbation form is classical [4]; the textbook treatment is Khalil's *Nonlinear Systems* [5, Ch.~11], with the composite-Lyapunov approach this derivation parallels due to Saberi and Khalil [6] and the control-theoretic corpus consolidated in Kokotović, Khalil and O'Reilly [7]. What is AAT-native is the statement of the stacking entirely in the sector-persistence template's (α, R, ρ) coordinates, so that every quantity in the stability condition is an already-named persistence quantity.

- **(S0) Regularity.** F_1, F_2 locally Lipschitz in all arguments; w_1, w_2 measurable (solutions in the Carathéodory sense).

Assumption (stacking premises S0–S4).

- **(S1) Lipschitz quasi-steady-state manifold.** There is $h : \mathcal{X}_2 \rightarrow \mathbb{R}^{n_1}$ with $F_1(h(x_2); x_2) = 0$ for all x_2 in the slow scope region, and $\|Dh(x_2)\| \leq L_h$.
- **(S2) Per-level sector conditions** — the template's (T2), per level, the fast one *uniformly parametrized* by the slow state:
 - Fast: $(x_1 - h(x_2))^T F_1(x_1; x_2) \geq \alpha_1 \|x_1 - h(x_2)\|^2$ for $\|x_1 - h(x_2)\| \leq R_1$, with α_1, R_1 uniform over the slow scope region.
 - Slow (reduced dynamics): with $\bar{F}_2(x_2) := F_2(x_2; h(x_2))$, there is x_2^* with $\bar{F}_2(x_2^*) = 0$ and $(x_2 - x_2^*)^T \bar{F}_2(x_2) \geq \alpha_2 \|x_2 - x_2^*\|^2$ for $\|x_2 - x_2^*\| \leq R_2$.
- **(S3) Bounded interconnection.** $\|F_2(x_2; x_1) - F_2(x_2; h(x_2))\| \leq L_{21} \|x_1 - h(x_2)\|$ on the joint scope region, and $M_2 := \sup_{\|x_2 - x_2^*\| \leq R_2} \|\bar{F}_2(x_2)\| < \infty$.
- **(S4) Bounded disturbance (Model D).** $\|w_1(t)\| \leq \rho_1, \|w_2(t)\| \leq \rho_2$, with standing reserve hypothesis $\Delta\rho_1^* := \alpha_1 R_1 - \rho_1 > 0$ (the fast level persists standalone with reserve to spare; if it does not, no timescale separation can rescue the stack).

Error coordinates $y_1 := x_1 - h(x_2), y_2 := x_2 - x_2^*$; joint scope region $\mathcal{B} := \{\|y_1\| \leq R_1\} \times \{\|y_2\| \leq R_2\}$.

Theorem. Under (S0)–(S4), define the slow-velocity bound $v_2^{\max} := M_2 + L_{21}R_1 + \rho_2$ and the separation threshold

Derived (two-timescale stacked persistence).

$$\epsilon_{\max} := \frac{\Delta\rho_1^*}{L_h v_2^{\max}} \quad (\text{M.2})$$

(with $\epsilon_{\max} = \infty$ when $L_h = 0$: a fixed target imposes no separation requirement). If

$$(\text{C1}) \quad \epsilon < \epsilon_{\max} \quad \text{and} \quad (\text{C2}) \quad \alpha_2 R_2 > \rho_2 + L_{21} R_1, \quad (\text{M.3})$$

then $\mathcal{B}$ is forward-invariant, and every trajectory starting in $\mathcal{B}$ satisfies

$$\limsup_{t \rightarrow \infty} \|y_1(t)\| \leq r_1 := \frac{\rho_1 + \epsilon L_h v_2^{\max}}{\alpha_1}, \quad \limsup_{t \rightarrow \infty} \|y_2(t)\| \leq r_2 := \frac{\rho_2 + L_{21} r_1}{\alpha_2}. \quad (\text{M.4})$$

Both levels persist: each ultimate-bound ball sits strictly inside its scope ball ((C1) gives $r_1 < R_1$; (C2) with $r_1 < R_1$ gives $r_2 < R_2$).

Derivation. (i) *A-priori slow velocity.* On $\mathcal{B}$, $\|F_2(x_2; x_1)\| \leq \|\bar{F}_2(x_2)\| + L_{21} \|y_1\| \leq M_2 + L_{21} R_1$ by (S3), so $\|\dot{x}_2\| \leq \epsilon v_2^{\max}$ — a scope-constant bound with no dependence on the ultimate bounds derived below (no circularity).

(ii) *The fast level as a template instance with manifold-drift disturbance.* Differentiating $y_1 = x_1 - h(x_2)$,

$$\dot{y}_1 = -F_1(y_1 + h(x_2); x_2) + \underbrace{w_1 - Dh(x_2) \dot{x}_2}_{w_1^{\text{eff}}}, \quad \|w_1^{\text{eff}}\| \leq \rho_1^{\text{eff}} := \rho_1 + \epsilon L_h v_2^{\max} \quad (\text{M.5})$$

by (S1) and (i). With $V_1 = \frac{1}{2} \|y_1\|^2$ and (S2-fast): $\dot{V}_1 \leq -\alpha_1 \|y_1\|^2 + \|y_1\| \rho_1^{\text{eff}}$ on $\mathcal{B}$. The slow level's motion enters the fast level as additional Model-D disturbance at rate $\epsilon L_h v_2^{\max}$ — the moving-target term — and the fast level's persistence condition $\alpha_1 R_1 > \rho_1^{\text{eff}}$ rearranges to (C1).

(iii) *The slow level as a template instance with renormalized disturbance.* $\dot{y}_2 = \epsilon[-\bar{F}_2(x_2) - (F_2(x_2; x_1) - \bar{F}_2(x_2)) + w_2]$, where the middle term is bounded by $L_{21}\|y_1\|$ by (S3). With $V_2 = \frac{1}{2}\|y_2\|^2$ and (S2-slow): $\dot{V}_2 \leq \epsilon[-\alpha_2\|y_2\|^2 + \|y_2\|(\rho_2 + L_{21}\|y_1\|)]$ on $\mathcal{B}$ — in slow time $\tau = \epsilon t$ the ϵ factors out, and the fast level, integrated out, leaves a residue equal to its settled bound times the interconnection sensitivity.

(iv) *Joint forward-invariance and the cascade bound.* On the face $\|y_1\| = R_1$ of $\partial\mathcal{B}$: $\dot{V}_1 \leq R_1(-\alpha_1 R_1 + \rho_1^{\text{eff}}) < 0$ by (C1). On the face $\|y_2\| = R_2$, with $\|y_1\| \leq R_1$ worst-case: $\dot{V}_2 \leq \epsilon R_2(-\alpha_2 R_2 + \rho_2 + L_{21}R_1) < 0$ by (C2). Each V_i decreases strictly (with uniform margin) on its own face, so a first-exit-time contradiction gives forward-invariance of $\mathcal{B}$, corners included, and all bounds above hold for all $t \geq 0$. Within $\mathcal{B}$ the fast inequality gives $\limsup \|y_1\| \leq r_1$ by the template's Grönwall argument; then for any $\delta > 0$ the slow effective disturbance is eventually below $\rho_2 + L_{21}(r_1 + \delta)$, giving $\limsup \|y_2\| \leq (\rho_2 + L_{21}(r_1 + \delta))/\alpha_2$ for every δ , hence the stated r_2 . $\square$

If the fast level is already settled at composition time ($\|y_1(0)\| \leq r_1 + \delta$ for some $\delta > 0$), the smaller ball $\{\|y_1\| \leq r_1 + \delta\}$ is itself forward-invariant (same face argument under (C1)), and (C2) weakens to

Derived (warm-start refinement — the cost of acting on transients).

$$(C2\text{-warm}) \quad \alpha_2 R_2 > \rho_2 + L_{21}(r_1 + \delta). \quad (M.6)$$

The gap between the cold-start and warm-start conditions, $L_{21}(R_1 - r_1)$ as $\delta \rightarrow 0$, is the quantitative price of engaging the slow level before the fast level has converged: early action does not void the guarantee — it raises the slow level's required reserve from settled-residue size to worst-case-transient size. 4.9's qualitative rule ("a slower process must not act before the faster process beneath it has converged") is thereby graded, not binary.

Remark (Tikhonov's prerequisite, within scope). Tikhonov-type results require the frozen fast dynamics to have a unique isolated equilibrium for each slow configuration. Under (S2-fast) this holds by construction within scope: if $F_1(y + h(x_2); x_2) = 0$ with $0 < \|y\| \leq R_1$, then $0 = y^\top F_1 \geq \alpha_1 \|y\|^2 > 0$ — contradiction — so $h(x_2)$ is the unique equilibrium in the radius- R_1 ball, and the sector additionally gives exponential frozen-dynamics attraction (more than Tikhonov asks; the derivation above in fact bypasses Tikhonov's theorem entirely via the composite-Lyapunov route). Outside the sector ball the guarantee is void: a fast level with multiple equilibria or limit cycles in its operating region — non-convex losses (K), cyclic strategic dynamics (16.1's R2 regime) — sits outside (S2-fast), and no amount of slowing the outer loop restores the guarantee. Slowing the slow level helps only against (C1) violations, never against a fast level that has no settled state to offer.

Derived (sector condition supplies the unique-root prerequisite).

Remark (relation to the template's stated preconditions). The template (G) states (T1)–(T3) for autonomous F and a constant disturbance bound. The two instances above are non-autonomous — the fast correction varies in time through $x_2(t)$, and the slow effective disturbance $\rho_2 + L_{21}\|y_1(t)\|$ is state-dependent and only *eventually* below the constant used — but the template's Lyapunov/Grönwall argument is pointwise in time and uses only the differential inequality, which holds uniformly, so the proofs transfer verbatim. This is a mild, stated extension of the template (uniform-in-time (T1)–(T2), eventually-bounded (T3)), not a bare instantiation.

Levels $k = 1, \dots, N$ (fastest to slowest) with rates $\epsilon_1 = 1 \gg \epsilon_2 \gg \dots \gg \epsilon_N$. Two structural premises beyond the per-level transcriptions of (S0)–(S4):

Derived (N-level stacking, nearest-neighbor coupling).

- **(S5) Nesting idealization.** Each level's quasi-steady-state manifold depends only on the level immediately above: $h_k = h_k(x_{k+1})$, Lipschitz L_{h_k} . (Dependence on all slower levels adds a bounded sum dominated by the adjacent term under the rate ordering; the nearest-neighbor statement keeps the constants legible.)

- **(S2-red) Reduced sector conditions for middle levels.** For $1 < k < N$ the sector condition is on the *reduced* dynamics with the level below substituted at its manifold, $\bar{F}_k(x_k; x_{k+1}) := F_k(x_k; h_{k-1}(x_k), x_{k+1})$ — note the manifolds compose, since h_{k-1} is itself a function of x_k — anchored at $h_k(x_{k+1})$, uniformly in x_{k+1} .

Each level k then carries three disturbance sources — its own ρ_k , the residue of the level below ($L_{k,k-1}\|y_{k-1}\|$), and the manifold drag from the level above ($(\epsilon_{k+1}/\epsilon_k)L_{h_k}v_{k+1}^{\max}$, with $v_{k+1}^{\max}$ the level- $(k+1)$ scope-constant velocity bound) — boundary terms zero at $k = 1$ and $k = N$. If for each k

$$(C-N) \quad \alpha_k R_k > \rho_k + L_{k,k-1} R_{k-1} + \frac{\epsilon_{k+1}}{\epsilon_k} L_{h_k} v_{k+1}^{\max}, \quad (M.7)$$

then $\prod_k \{\|y_k\| \leq R_k\}$ is forward-invariant by the same per-face argument, and the cascade limsup propagates upward level by level, giving the renormalization recursion on ultimate bounds

$$r_k = \frac{\rho_k + L_{k,k-1} r_{k-1} + (\epsilon_{k+1}/\epsilon_k) L_{h_k} v_{k+1}^{\max}}{\alpha_k}, \quad r_0 := 0. \quad (M.8)$$

The recursion is well-founded in both directions: from-above terms use a-priori scope constants, from-below terms use already-derived r_{k-1} . For general (non-nearest-neighbor) coupling, the per-face argument is replaced by the standard vector-Lyapunov / M-matrix condition for interconnected systems [5, §9.5] — imported machinery, named here as the general gate without re-derivation.

With Wiener-process disturbance the manifold subtraction acquires an Itô correction requiring (S1') $h \in C^2$ with bounded Hessian, and the drift-disturbance bookkeeping above becomes a joint exit-time problem across coupled levels. The expected shape — mean-square ultimate bounds with the template's $1/\sqrt{\alpha}$ scaling and the region-aware stopped form of A Prop A.1S, plus a curvature term in ρ^{eff} — is genuinely additional work, not a transcription, and is *not claimed here*. The Model D scope is the honest match for the stacking's primary AAT use cases (structural adaptation, consolidation scheduling, macro-clocks), where slow-level motion is drift-like.

Formulation (Model S extension — conditional, not yet derived).

Epistemic Status.

Exact for the two-timescale theorem, the warm-start refinement, and the N -level nearest-neighbor form, under the named premises (S0)–(S5)/(C-N) — derived by elementary Lyapunov arguments stacked on the sector-persistence template, independently re-derived and verified in full (spike trail 2026-06-10). The general-coupling N -level gate is *imported* (vector-Lyapunov / M-matrix, [5, §9.5]) and the Model S extension is *open* — both stated as such above, neither carried by the exact label.

What the promotion does and does not close. The predecessor sketch's open problem — AAT has trigger conditions but no dynamics for structural adaptation — is *conditionalized, not closed*. The theorem covers any deeper-level dynamics satisfying (S1)–(S5); whether AAT's actual deeper levels satisfy them is exactly the question that remains open, and it is a real one: structural adaptation is plausibly discrete (adding a model component, restructuring a hierarchy), making $\dot{x}_2$ a jump process for which (S0)'s Carathéodory framing and (S1)'s differentiable manifold do not hold as stated. The claim-type has changed from “AAT's adaptive hierarchy is stable” (never derivable as stated) to “a premise-conditional stacking theorem with named scope” — the strengthening is real (closed-form $\epsilon_{\max}$, graded warm/cold gap, unique-root prerequisite supplied rather than assumed), and so is the remaining gap.

The load-bearing premise is (S2-fast)'s *uniformity*, not the sector inequality itself. Because (α_1, R_1) are uniform over the slow scope region and h absorbs all equilibrium motion, the entire interconnection burden lands on the single drift term $\epsilon L_h v_2^{\max}$; a per-frozen- x_2 sector with non-uniform constants would not suffice. This is the standard strength

assumption of the singular-perturbation literature (it is essentially Saberi-Khalil's interconnection-condition pattern), and it is precisely where realistic fast levels — non-convex, multi-equilibrium — exit scope.

Max attainable within current premises: exact (attained for Model D). Extensions that would widen scope rather than strengthen the claim: Model S (named above), jump-process slow dynamics, non-Euclidean metrics via AI's (CT2).

Discussion.

The convergence constraint is now a derived threshold, and its two failure modes are dual. 4.9's $\nu_{n+1} \ll \nu_n$ enters the theorem as (C1) with the closed form $\epsilon_{\max} = \Delta\rho_1^*/(L_h v_2^{\max})$: the faster level's *spare adaptive reserve* — a quantity the template already names — divided by the rate at which the slower level drags the faster level's target. The two classic nesting pathologies separate cleanly into the theorem's two conditions. *Micromanagement* is a (C1) violation: the slow level acts at fast tempo, the moving-target disturbance $\epsilon L_h v_2^{\max}$ exhausts the fast level's reserve, and the fast level thrashes — and the damage then propagates back upward, because the fast level's inflated residue $L_{21}r_1$ raises the slow level's effective disturbance in turn. *Catastrophic forgetting* is a (C2) violation: the slow (consolidated) level is too sensitive to fast-level transients — L_{21} too large — and the fast loop's activity overwrites it even under perfect timescale separation. The two pathologies several auditors read as “the same failure” are precisely the two conditions of one theorem, violated separately.

Timescale separation is trust, quantified. “Each level has a stable attractor given the levels above it” reads as mutual trust within an architecture: the fast layer trusts the slow layer not to pull the rug out mid-convergence ((C1) — the manager's rate of goal-change stays under $\epsilon_{\max}$), and the slow layer trusts the fast layer to settle so there is a steady output to act on ((C2) — the worker's residue stays within the manager's tolerance $L_{21}r_1$). The warm-start refinement prices the violation: a manager who engages before the worker converges needs $L_{21}(R_1 - r_1)$ of extra reserve — the difference between reacting to settled output and reacting to worst-case transients.

Integrating out a level is a renormalization step; the template is the fixed-point form. Step (iii)'s structure — the fast level, integrated out, maps the slow level onto the *same* template form with renormalized disturbance $\rho_{k+1} \mapsto \rho_{k+1} + L_{k+1,k} r_k$ — is a renormalization-group step in the sense in which Chen, Goldenfeld and Oono showed RG and singular-perturbation asymptotics to be one method [8]. The sector-persistence template is form-invariant under this coarse-graining: every level of the hierarchy presents the identical $(\alpha, R, \rho^{\text{eff}})$ persistence problem with parameters renormalized by its neighbors. This is the precise sense in which AAT's persistence architecture is scale-free, and it is the lineage of Haken's slaving principle — fast modes enslaved by slow order parameters [9] — restated in persistence coordinates.

Renormalizing generative models are the constructive discrete-time instance. Friston and colleagues' scale-free active inference [10] builds generative models by recursive blocking transformations over state space and time, such that “higher levels encode sequences of sequences” and belief updating slows as one ascends levels — i.e., the timescale separation is enforced *architecturally* (one slow tick per block of K fast ticks, $\epsilon_k/\epsilon_{k+1} = 1/K$ by construction) rather than assumed, with per-level convergence inherited from coordinate descent. RGM thus supplies what (S1)–(S2) ask for, by construction, in the discrete categorical setting — an existence proof that the premises are realizable, not a competing stability argument (the paper proves no composite-stability theorem and does not claim to). The same architecturally-enforced separation appears in AAT's own composition machinery as 13.1's macro-clock $K_c \gg 1$ and #def-auxilia-hierarchy's (H5).

The LLM adaptive stack is the N -level hierarchy, unlabeled. Pretraining $\rightarrow$ fine-tuning $\rightarrow$ LoRA-style adaptation $\rightarrow$ in-context learning $\rightarrow$ retrieval state $\rightarrow$ within-generation dynamics maps onto the timescale ladder with each mechanism orders of magnitude faster and more transient than the one below. The framework's stability requirement needs no clean “parametric vs structural” labels — only that adjacent rates satisfy (C-N) — and it explains the practitioner's instinct against, e.g., updating base weights while a fast retrieval loop is mid-flight: that is a (C1) violation at the weights/retrieval boundary. The catastrophic-forgetting literature lives at the same boundary as a (C2) violation, per the dual reading above.

Strategy’s own timescale. 9.10 places strategic mismatch between fast epistemic updates and slow objective revision; the stacking theorem is the formal backing for treating those as separately-stable levels — with the caveat that the strategic level’s $\alpha_\Sigma = 1/(n + 1)$ decay makes (S2)’s constant- α premise hold only under experience discounting, exactly as that segment already requires.

Working Notes.

- **Open (which AAT mechanisms satisfy the premises).** The honest gap left by the promotion: structural adaptation as actually triggered by 4.8 is plausibly a *jump process* (discrete model-class moves), outside (S0)/(S1) as stated. Two follow-up directions: an impulsive/hybrid-systems extension of the stacking (the slow level as a jump process on the slow manifold), and formalizing the RGM correspondence ([10]) as the discrete-time constructive instance where the premises hold by architecture. Spike trail: spikes/spike-multi-timescale-stacking-2026-06-10.md.
 - **Open (Model S stacking).** Named conditional above; the joint exit-time bookkeeping across coupled levels is the work item. Prop A.1S’s region-aware form is the template to extend.
 - **Open (non-Euclidean metrics).** The stacking is stated in the Euclidean template; restating with AIs (CT1)–(CT3) per level would cover Fisher-metric fast levels (statistical-manifold learners) — likely routine but unchecked.
 - **Belongs elsewhere (ELI internal coherence).** The trust reading extends to ELI architecture: a narrative “I” that constantly interrupts its own sub-agents violates (C1) at the interiority boundary (“paralyzed by frantic internal micromanagement” — Gemini, AUDIT-WORKING-193847). Carried at #def-auxilia-hierarchy (H5) and its temporal-sovereignty discussion; not this segment’s claim to make.
 - Incidental audit gold from the 2026-05-31 lift is integrated: the trust/manager-worker framing, the LLM-stack mapping, the micromanagement/forgetting unification (sharpened here into the dual (C1)/(C2) reading), and the “does slowing the outer loop guarantee convergence?” reader question (answered in the Tikhonov remark: no — a multi-equilibrium fast loop is outside (S2-fast) regardless of ϵ) all landed in the body above. Provenance: the pre-promotion segment history and the spike trail.
-

Appendix N

Sketch: Does the Model-S Containment Dichotomy License Structural-Adaptation Genericity?

The Model-S segments claim Corollary A.1S.1 “sharpens the hand-off into #result-structural-adaptation-necessity — the structural-adaptation trigger is *generic, not exceptional*, for a sufficiently long-lived agent.” This sketch isolates that hand-off as an open question: as stated it is plausibly a mild overclaim, and the honest move is to *derive the bridge* rather than soften the prose.

The Open Question.

#deriv-sector-condition Corollary A.1S.1 establishes (exactly) that under additive stochastic forcing the persistence-region first-exit probability is $P(\tau_R < \infty) = 1$ — a long-lived agent leaves the sector-condition region $\mathcal{B}_R$ with certainty. The Model-S segments and #deriv-stochastic-non-exit then assert this “sharpens the hand-off” so that structural adaptation is *generic, not exceptional*.

Discussion.

But #result-structural-adaptation-necessity derives necessity from a **different trigger**: model-class inadequacy ($\mathcal{F}(\mathcal{M}) < 1 - \varepsilon$ — no model in the current class can represent reality). It does not reference Model-S, region-exit, or Corollary A.1S.1; its mechanism is fitness-driven, not noise-driven. And the corollary’s own explicit instance (the OU benchmark) is **positively recurrent**: the process exits any bounded region a.s. *and returns* a.s. — exit-and-return, not escape.

So there is a gap between what Cor A.1S.1 proves and what the hand-off prose claims:

- **What the corollary licenses (honest narrow claim):** the *pathwise-forever parametric-containment guarantee* is structurally unavailable under stochastic forcing; a long-lived agent generically enters the regime where the parametric persistence guarantee is no longer *proven* (the out-of-region regime #result-structural-adaptation-necessity addresses).
- **What it does not license without further argument:** that *model-class* structural adaptation is generically *necessary*. Recurrent region-exit (exit-and-return) is not model-class inadequacy; a transient noise excursion beyond $\mathcal{B}_R$ that the (un-posed-but-possibly-still-inward) correction pulls back is not a structural-adaptation trigger in the #result-structural-adaptation-necessity sense.

Proposed Direction (rigor pending).

Per strengthen-before-soften, the preferred resolution is **derive the bridge**, not tighten the prose. Candidate conditions under which certain region-exit *would* imply structural-adaptation necessity, to be attempted:

Discussion.

1. **A2' genuinely fails outside $\mathcal{B}_R$** (not merely “not posited”): if the correction function provably ceases to point inward beyond R for the agent’s model class, then a.s. exit composes with non-recovery, and recurrent exit *does* force the structural regime. This connects to the sub-scope β / model-class-capacity reading of R in #form-sector-condition.
2. **Excursion statistics interact with $\mathcal{F}(\mathcal{M})$** : if recurrent large excursions degrade *effective* model-class fitness (the model fitted within $\mathcal{M}$ becomes systematically wrong during excursions faster than it re-converges between them), the noise-driven and fitness-driven triggers couple, and genericity follows from the coupling rather than from exit alone.
3. **Timescale-separation argument**: if structural adaptation operates so much slower than parametric correction that the *fraction of time* spent out-of-region (positive under recurrence, however small) accumulates an unbounded structural-adaptation debt over an unbounded horizon.

Each is a derivation target; any one that closes turns the hand-off from asserted to proven and is a genuine strengthening (a new conditional result, not a softening). If all provably fail, the honest outcome is the no-go protocol applied to the hand-off claim (per doc/audit-routing-instructions.md §4) and the Model-S Discussion/Findings prose tightened to the narrow licensed claim above — but only *after* the strengthening attempt is exhausted and recorded.

Epistemic Status.

Sketch. Direction identified, formalization not attempted. This segment makes **no** claim that the bridge holds or fails; it records that the hand-off currently asserted in #deriv-sector-condition / #deriv-stochastic-non-exit Discussion/Findings is **plausibly a mild overclaim at the Discussion/Findings tier** (a Gate-2 class concern: a claim that *sounds* like it follows from Cor A.1S.1 but is not derived from it), and that the resolution path is to attempt the bridge derivation. The narrow licensed claim (pathwise-parametric-guarantee unavailability $\Rightarrow$ generic entry into the out-of-region *regime*) is sound and uncontroversial; the disputed increment is “ $\Rightarrow$ model-class structural adaptation is generically *necessary*.” Ceiling: a closed bridge would be derived (conditional) on the named condition; a clean impossibility would be a no-go result. Until attempted, this stays sketch.

Discussion.

Why this matters enough to surface rather than leave as a soft ledger row: the genericity claim is load-bearing downstream — it is the form in which the Model-S no-go’s consequence reaches the ELI / long-lived-agent persistence framing (“in any genuinely stochastic environment, structural adaptation is not an edge case”). If the bridge holds under a named condition, that is a sharp, citable strengthening of exactly the kind the framework’s no-go signature tends to produce. If it does not, the framework’s honesty requires the hand-off prose to state only the narrow licensed claim. Either way the question should be visible in the argument’s structure, not invisible in a tracking file — which is the point of carrying it as an exploratory OUTLINE node.

Working Notes.

- Provenance: parent-generated 2026-05-16 during the audit-routing cycle, from a first-hand deep-read of #result-structural-adaptation-necessity against the Model-S segments (not a cluster audit finding). Originally filed as consolidated-ledger

research-seed **S30**; Joseph adjudicated it up to a first-class exploratory OUTLINE node (“it will get its day in the light”) rather than a soft ledger row. The ledger row is correspondingly demoted to a pointer here.

- Segment-vs-spike judgment (recorded so it is not re-litigated): this is recognized open *territory* between two landed results with **no attempt yet** — segment-shaped (visible structural node), not spike-shaped (a spike is the reasoning trail of an *attempt*). The spike is the eventual *work*: when this is picked up, a strengthening spike attempts directions 1–3; math lands back here (or in a new derivation/result segment), per math-lives-in-segments.
 - Related: the #disc-identifiability-floor family — if the bridge is a no-go, the “no horizon-independent non-exit certificate” structural absence may itself compose with a structural-adaptation-necessity floor; worth checking whether this is an identifiability-floor-shaped instance.
-

Appendix O

Derivation: Discrete-Time Sector Condition

Discrete-time analogs of Props A.1, A.1S, and A.2 via contraction mapping, closing the fluid-limit gap (GA-5) between the event-driven dynamics (3.1) and the continuous-time Lyapunov results in A. Three results land. **Prop DA.1 (Model D)**: under deterministic bounded disturbance, the discrete steady-state mismatch equals the continuous steady-state exactly — zero fluid-limit gap. **Prop DA.1S (Model S)**: under stochastic disturbance, the discrete and continuous steady-state variances differ by an additive correction of order $\eta^* c_{\max}^2 / c_{\min}^2$ — quantitatively small when this conditioning-ratio-over-event-rate product is small. **Prop DA.2 (adaptive reserve)**: the discrete adaptive reserve matches the continuous form modulo the same correction.

A critical structural observation underwrites all three: the discrete sector condition (DA2') needs a *strictly stronger Lipschitz upper bound* on the correction function — strictly stronger than an inner-product upper bound — because the discrete recurrence's quadratic term $(\eta^*)^2 \|F_d\|^2$ requires the norm bound, not merely the projection. This is satisfied for Bayesian updates (the posterior lies between prior and data) and for gradient descent on smooth losses (where the Lipschitz constant is the smoothness parameter); pathological correction functions with large transverse components can satisfy the inner-product bound while violating the norm bound. Together with the no-overshoot step-size condition $\eta^* < 2c_{\min}/c_{\max}^2$, the fluid limit is formally justified — the continuous-time analysis of A is the limit of the discrete results here, not an idealization.

Formal Expression.

Setup. The discrete mismatch dynamics at event step k are:

Definition (Discrete Dynamics).

$$\delta_{k+1} = \delta_k - \eta^* F_d(\delta_k) + w_k \tag{O.1}$$

where η^* is the update gain (3.6), F_d is the discrete correction direction, and w_k is the per-step disturbance. The continuous correction function $F(\mathcal{J}, \delta)$ from A decomposes as $F = \nu \cdot \eta^* \cdot F_d$ at event rate ν .

(DA2') Discrete Sector-Lipschitz Condition.

There exist constants $c_{\min} > 0$ and $c_{\max} < 2/\eta^*$ such that for all $\|\delta\| \leq R$:

Assumption DA2' (discrete-sector-condition).

(DA2'squa) Lower sector bound (directional fidelity):

$$\delta^T F_d(\delta) \geq c_{\min} \|\delta\|^2 \quad (\text{O.2})$$

(DA2'squob) Lipschitz bound (bounded correction magnitude):

$$\|F_d(\delta)\| \leq c_{\max} \|\delta\| \quad (\text{O.3})$$

The **lower bound** (DA2'a) is directional fidelity — the correction points inward, identical to the continuous sector condition (A2') from 4.2 via 4.3.

The **Lipschitz bound** (DA2'b) controls the *magnitude* of the correction, not merely its projection onto the mismatch direction. The combined constraint $c_{\max} < 2/\eta^*$ is the **no-overshoot condition**: each correction step must not reverse the mismatch. This is the classical step-size condition $\eta^* < 2/L$ for gradient descent (where L is the Lipschitz constant of the gradient). For Bayesian updates, this is satisfied by construction — the posterior lies between prior and data.

Why DA2'b is stronger than an inner-product upper bound. A two-sided inner-product condition $\delta^T F_d(\delta) \leq c_{\max} \|\delta\|^2$ constrains only the projection of F_d onto δ . By Cauchy-Schwarz, the Lipschitz bound (DA2'b) implies the inner-product upper bound: $\delta^T F_d(\delta) \leq \|\delta\| \cdot \|F_d(\delta)\| \leq c_{\max} \|\delta\|^2$. But the converse fails — a correction function with a large transverse component (orthogonal to δ) can satisfy the inner-product bound while violating the norm bound. The proofs below (especially DA.1S) require the norm bound $\|F_d(\delta)\|^2 \leq c_{\max}^2 \|\delta\|^2$, which follows from DA2'b but not from an inner-product condition alone.

Scalar case. In one dimension, DA2'a and DA2'b together reduce to the classical sector condition $c_{\min} \leq F_d(\delta)/\delta \leq c_{\max}$, since norm and inner product coincide. No generality is lost for scalar systems.

Relationship to the continuous-time condition. The continuous sector condition (A2'/GA-3) is a one-sided inner-product bound $\delta^T F \geq \alpha \|\delta\|^2$ — this suffices for continuous-time Lyapunov analysis because $\dot{V}$ involves only $\delta^T F$, not $\|F\|$. Discretization introduces the quadratic term $(\eta^*)^2 \|F_d\|^2$ (see DA.1S proof), which requires the Lipschitz bound. This is the standard sector-vs-Lipschitz distinction in nonlinear systems theory.

Contraction factor. Under DA2', the per-step Lyapunov function $V_k = \frac{1}{2} \|\delta_k\|^2$ satisfies (in the zero-disturbance case $w_k = 0$):

Derived (contraction, from DA2').

$$\|\delta_{k+1}\|^2 = \|\delta_k - \eta^* F_d(\delta_k)\|^2 = \|\delta_k\|^2 - 2\eta^* \delta_k^T F_d(\delta_k) + (\eta^*)^2 \|F_d(\delta_k)\|^2 \quad (\text{O.4})$$

Applying DA2'a (lower sector bound on $\delta^T F_d$) and DA2'b (Lipschitz bound on $\|F_d\|$):

$$\|\delta_{k+1}\|^2 \leq (1 - 2\eta^* c_{\min} + (\eta^*)^2 c_{\max}^2) \|\delta_k\|^2 = \lambda_{\text{eff}}^2 \|\delta_k\|^2 \quad (\text{O.5})$$

where:

$$\lambda_{\text{eff}}^2 = 1 - 2\eta^* c_{\min} + (\eta^*)^2 c_{\max}^2 \quad (\text{O.6})$$

Stability condition. $\lambda_{\text{eff}}^2 < 1$ requires $2\eta^* c_{\min} > (\eta^*)^2 c_{\max}^2$, i.e.:

$$\eta^* < \frac{2c_{\min}}{c_{\max}^2} \quad (\text{O.7})$$

This is automatically satisfied when $c_{\min} \approx c_{\max}$ (well-conditioned correction), recovering the standard step-size condition $\eta^* < 2/c_{\max}$. For ill-conditioned corrections ($c_{\min} \ll c_{\max}$), the constraint is tighter. For Bayesian updates with bounded condition number, both conditions are satisfied.

Scalar (colinear) specialization. When $F_d(\delta) \parallel \delta$ (scalar system or colinear correction), $\|F_d(\delta)\| = |F_d(\delta)/\delta| \cdot \|\delta\|$ and the contraction factor simplifies to $\lambda = \max(|1 - \eta^* c_{\min}|, |1 - \eta^* c_{\max}|)$, the classical form. The general vector formula λ_{eff}^2 reduces to λ^2 in this case.

With disturbance $w_k \neq 0$:

$$\|\delta_{k+1}\|^2 \leq \lambda_{\text{eff}}^2 \|\delta_k\|^2 + 2\|\delta_k\| \|w_k\| + \|w_k\|^2 \quad (\text{O.8})$$

Proposition DA.1: Bounded Mismatch (Deterministic). **Statement.** Under DA2' with $\eta^* < 2c_{\min}/c_{\max}^2$ and bounded per-step disturbance $\|w_k\| \leq \rho_{\text{step}}$, the mismatch is ultimately bounded:

Derived (DA.1, discrete bounded mismatch).

$$R_D^* = \frac{\rho_{\text{step}}}{1 - \lambda_{\text{eff}}} \quad (\text{O.9})$$

Proof. By the triangle inequality: $\|\delta_{k+1}\| = \|(\delta_k - \eta^* F_d(\delta_k)) + w_k\| \leq \|\delta_k - \eta^* F_d(\delta_k)\| + \|w_k\|$. The contraction bound gives $\|\delta_k - \eta^* F_d(\delta_k)\| \leq \lambda_{\text{eff}} \|\delta_k\|$. Therefore:

$$\|\delta_{k+1}\| \leq \lambda_{\text{eff}} \|\delta_k\| + \rho_{\text{step}} \quad (\text{O.10})$$

This is an affine contraction with $\lambda_{\text{eff}} < 1$. By the Banach fixed-point theorem, all trajectories starting in $\mathcal{B}_R$ converge to the ball of radius $R_D^* = \rho_{\text{step}}/(1 - \lambda_{\text{eff}})$, provided $R_D^* < R$. $\square$

Recovery of continuous result. In the fluid limit ($\eta^* \rightarrow 0, \nu \rightarrow \infty, \nu\eta^* = \mathcal{J}$ fixed), $\lambda_{\text{eff}}^2 = 1 - 2\eta^* c_{\min} + O((\eta^*)^2)$, so $\lambda_{\text{eff}} \rightarrow 1 - \eta^* c_{\min}$ and $\rho_{\text{step}} \rightarrow \rho/\nu$. Then:

$$R_D^* = \frac{\rho/\nu}{\eta^* c_{\min}} = \frac{\rho}{\nu\eta^* c_{\min}} = \frac{\rho}{\alpha} \quad (\text{O.11})$$

recovering Prop A.1 exactly. The discrete-to-continuous gap for Model D steady state is **zero**.

Proposition DA.2: Adaptive Reserve (Discrete). **Statement.** Under the conditions of DA.1, the agent can absorb an additional per-step disturbance of:

Derived (DA.2, discrete adaptive reserve).

$$\Delta \rho_{\text{step}}^* = (1 - \lambda_{\text{eff}})R - \rho_{\text{step}} \quad (\text{O.12})$$

Proof. Identical to Prop A.2: the new $R_D^* = (\rho_{\text{step}} + \Delta \rho)/(1 - \lambda_{\text{eff}})$ must satisfy $R_D^* \leq R$. $\square$

The structure is identical to the continuous adaptive reserve $\Delta \rho^* = \alpha R - \rho$. The per-event contraction rate is $(1 - \lambda_{\text{eff}})$; the per-unit-time rate is $\nu(1 - \lambda_{\text{eff}})$. In the fluid limit, $\nu(1 - \lambda_{\text{eff}}) \rightarrow \nu \cdot \eta^* c_{\min} = \alpha$, recovering the continuous result. (Note: the per-event rate $(1 - \lambda_{\text{eff}})$ converges to $\eta^* c_{\min}$, not to α directly — the factor of ν converts between per-event and per-unit-time.)

Proposition DA.1S: Stochastic Bounded Mismatch (Discrete). **Statement.** Under DA2' with $\eta^* < 2c_{\min}/c_{\max}^2$ and i.i.d. zero-mean disturbance $\mathbb{E}[w_k] = 0$, $\mathbb{E}[\|w_k\|^2] = \sigma_{\text{step}}^2$, the mismatch satisfies:

Derived (DA.1S, discrete stochastic bounded mismatch).

$$\mathbb{E}[\|\delta_k\|^2] \leq \lambda_{\text{eff}}^{2k} \|\delta_0\|^2 + \frac{\sigma_{\text{step}}^2}{1 - \lambda_{\text{eff}}^2} \quad (\text{O.13})$$

where $\lambda_{\text{eff}}^2 = 1 - 2\eta^* c_{\min} + (\eta^*)^2 c_{\max}^2$.

Proof. Define $V_k = \|\delta_k\|^2$. Taking expectations of the squared update:

$$\mathbb{E}[V_{k+1} \mid \delta_k] = \|\delta_k - \eta^* F_d(\delta_k)\|^2 + \sigma_{\text{step}}^2 \quad (\text{O.14})$$

The first term expands as:

$$\|\delta_k - \eta^* F_d(\delta_k)\|^2 = V_k - 2\eta^* \delta_k^T F_d(\delta_k) + (\eta^*)^2 \|F_d(\delta_k)\|^2 \quad (\text{O.15})$$

By DA2'a (lower sector bound): $\delta_k^T F_d(\delta_k) \geq c_{\min} V_k$.

By DA2'b (Lipschitz bound): $\|F_d(\delta_k)\|^2 \leq c_{\max}^2 V_k$.

Note that the second step requires the *norm* bound DA2'b, not merely an inner-product upper bound — this is where the Lipschitz condition is essential. Combining:

$$\mathbb{E}[V_{k+1} \mid \delta_k] \leq \lambda_{\text{eff}}^2 V_k + \sigma_{\text{step}}^2 \quad (\text{O.16})$$

The condition $\eta^* < 2c_{\min}/c_{\max}^2$ ensures $\lambda_{\text{eff}}^2 < 1$. This is a supermartingale (when V_k is large enough). Iterating:

$$\mathbb{E}[V_k] \leq \lambda_{\text{eff}}^{2k} V_0 + \frac{\sigma_{\text{step}}^2}{1 - \lambda_{\text{eff}}^2} \quad (\text{O.17})$$

The steady-state mean-square mismatch is $\sigma_{\text{step}}^2/(1 - \lambda_{\text{eff}}^2)$. $\square$

Recovery of continuous result. In the fluid limit: $\sigma_{\text{step}}^2 \rightarrow n\sigma_w^2/\nu$, $\lambda_{\text{eff}}^2 \rightarrow 1 - 2\eta^*c_{\min}$, and $(1 - \lambda_{\text{eff}}^2) \rightarrow 2\eta^*c_{\min}$. The steady-state becomes $n\sigma_w^2/(2\nu\eta^*c_{\min}) = n\sigma_w^2/(2\alpha)$, recovering Prop A.1S exactly.

The discrete-to-continuous gap for Model S variance is $O(\eta^*c_{\max}^2/c_{\min}^2) = O(c_{\max}^2/(c_{\min}^2\nu))$, dominated by the conditioning ratio $c_{\max}^2/c_{\min}^2$. Substituting $\sigma_{\text{step}}^2 = n\sigma_w^2/\nu$ and $\eta^* = \mathcal{J}/\nu$ into $V_{\text{ss}} = \sigma_{\text{step}}^2/(1 - \lambda_{\text{eff}}^2)$ and Taylor-expanding at small η^* :

$$\frac{V_{\text{ss}}}{V_c} = \frac{1}{1 - \eta^*c_{\max}^2/(2c_{\min})} = 1 + \frac{\eta^*c_{\max}^2}{2c_{\min}} + O((\eta^*)^2) \quad (\text{O.18})$$

so

$$V_{\text{ss}} - V_c \approx \frac{n\sigma_w^2 c_{\max}^2}{4c_{\min}^2 \nu} \quad (\text{O.19})$$

— a leading correction that scales as $1/\nu$. The $(\eta^*)^2\|F_d\|^2$ term in the per-step recurrence enters $1 - \lambda_{\text{eff}}^2$ at order $(\eta^*)^2$, but the steady-state ratio $\sigma_{\text{step}}^2/(1 - \lambda_{\text{eff}}^2)$ inverts the leading $2\eta^*c_{\min}$ contraction, yielding $O(\eta^*)$ asymptotic gap rather than $O((\eta^*)^2)$. This quantifies the error introduced by GA-5 and confirms it is small whenever $c_{\max}^2/(c_{\min}^2\nu) \ll 1$.

Fluid Limit Theorem.

Statement. Let F_d be Lipschitz continuous with constant L_F on $\mathcal{B}_R$. Let $\delta^{(\nu)}(t)$ denote the piecewise-constant interpolation of the discrete trajectory at event rate ν , and $\delta(t)$ the solution of the continuous ODE $d\delta/dt = -F(\mathcal{J}, \delta) + w(t)$ with $F = \nu\eta^*F_d$. Then:

Derived (Conditional on Lipschitz regularity).

$$\sup_{t \in [0, T]} \|\delta^{(\nu)}(t) - \delta(t)\| \leq C \cdot \frac{\eta^*c_{\max}}{\nu^{1/2}} \quad (\text{O.20})$$

for a constant C depending on T, L_F , and R .

Sketch. This follows from the standard ODE-approximation theory for Euler schemes (Kushner & Yin, 2003, Ch. 5). The discrete update $\delta_{k+1} = \delta_k - \eta^*F_d(\delta_k) + w_k$ is a forward Euler step for the ODE with step size $1/\nu$. The Lipschitz condition on F_d ensures the local truncation error is $O(1/\nu^2)$. Summing over $O(\nu T)$ steps and applying the Gronwall inequality gives the $O(1/\nu^{1/2})$ bound.

For Model D (deterministic): the steady-state gap is exactly zero (both discrete and continuous converge to the same fixed point). The fluid-limit error affects only transients.

For Model S (stochastic): the steady-state variance gap is $O(\eta^*c_{\max}^2/c_{\min}^2) = O(c_{\max}^2/(c_{\min}^2\nu))$, dominated by the conditioning ratio $c_{\max}^2/c_{\min}^2$. The $(\eta^*)^2\|F_d\|^2$ term in the per-step recurrence enters $1 - \lambda_{\text{eff}}^2$ at order $(\eta^*)^2$, but inverting the leading $2\eta^*c_{\min}$ in $V_{\text{ss}} = \sigma_{\text{step}}^2/(1 - \lambda_{\text{eff}}^2)$ produces an $O(\eta^*)$ asymptotic gap, not $O((\eta^*)^2)$ — see the recovery calculation in Prop DA.1S above.

Epistemic Status.

DA.1, DA.1S, and DA.2 are *conditional* on DA2', which now includes both a sector lower bound (DA2'a) and a Lipschitz upper bound (DA2'b). The Lipschitz condition is strictly stronger than an inner-product upper bound — this is a real assumption, not a technicality. For Bayesian updates and gradient descent on smooth losses, DA2'b is satisfied (the correction magnitude is controlled by the loss curvature). For agents with non-smooth or pathological correction functions, DA2'b must be verified independently.

The proofs themselves are standard contraction-mapping (DA.1/DA.2) and supermartingale (DA.1S) arguments. The step-size constraint $\eta^* < 2c_{\min}/c_{\max}^2$ is essential and is stated explicitly in each proposition.

Fluid limit theorem is *conditional* on Lipschitz regularity of F_d — a standard regularity condition satisfied by all correction functions in the verified instances table (4.3). The convergence rate follows from classical ODE approximation theory (Kushner & Yin, 2003); the application to the AAT mismatch dynamics is new but the mathematics is not.

Max attainable: conditional — the Lipschitz bound DA2'b is inherent to the discrete-time treatment and cannot be replaced by a weaker inner-product condition without losing the norm bound used in the proofs.

Discussion.

GA-5 is closed. The fluid-limit bridging assumption is no longer required as an ungrounded assumption. For Model D, the discrete and continuous steady states are identical — the gap is zero. For Model S, the gap is $O(\eta^* c_{\max}^2 / c_{\min}^2) = O(c_{\max}^2 / (c_{\min}^2 \nu))$ in variance, quantitatively bounded and small in the regime where $c_{\max}^2 / (c_{\min}^2 \nu) \ll 1$. The continuous-time results in A are formally justified as the fluid limit of the discrete results here.

DA2' is a sector-Lipschitz condition, not a pure sector condition. The continuous-time treatment (A) requires only the one-sided sector bound $\delta^T F \geq \alpha \|\delta\|^2$ because the Lyapunov derivative $\dot{V}$ depends only on $\delta^T F$. Discretization introduces a quadratic term $(\eta^*)^2 \|F_d\|^2$ that requires the Lipschitz bound $\|F_d\| \leq c_{\max} \|\delta\|$. This is the standard distinction between sector conditions and Lipschitz conditions in nonlinear systems theory. For all agents in the verified instances table (4.3), the Lipschitz bound holds alongside the sector bound, so DA2' imposes no additional restriction in practice.

The step-size constraint has two forms. For well-conditioned corrections ($c_{\min} \approx c_{\max}$), the constraint reduces to $\eta^* < 2/c_{\max}$ — the classical form. For ill-conditioned corrections ($c_{\min} \ll c_{\max}$), the tighter constraint $\eta^* < 2c_{\min}/c_{\max}^2$ applies. This distinction is invisible in scalar systems but matters for high-dimensional agents with anisotropic correction.

No downstream results change qualitatively. The persistence condition, adaptive reserve, and adversarial scaling laws derived in A and 15.2 hold as stated. The discrete framework provides sharper constants (replacing α with $\nu(1 - \lambda_{\text{eff}})$) and a quantitative transient error bound, but the qualitative structure is unchanged.

Part I's formal chain is now complete. The prediction chain:

$$\text{gain principle + B1} \xrightarrow{\text{derived}} \text{sector condition} \xrightarrow{\text{Lyapunov/contraction}} \text{persistence, reserve, scaling} \quad (\text{O.21})$$

holds in both continuous time (via A) and discrete time (via this segment), with a quantitative bridge between them (the fluid limit theorem).

Working Notes.

- Non-stationary gain: when η^* varies over time (as in Kalman filtering with growing P^-), the contraction factor λ changes per step. The ultimate bound still holds if $\sup_k |\lambda_k| < 1$, but the transient analysis requires time-varying contraction arguments.

- Heavy-tailed disturbances: DA.1S assumes finite second moment. Sub-Weibull or heavy-tailed w_k would need truncation arguments or alternative Lyapunov functions. These extreme cases are better modeled as triggers for structural adaptation (4.8).
- Non-i.i.d. disturbances: correlated w_k (e.g., Markov environment noise) weakens the supermartingale argument. The contraction still holds per-step, but the steady-state bound requires mixing conditions. This connects to the adversarial coupling analysis in 15.2.

Incidental audit gold (gold-lift, 2026-05-31). Cross-audit ideation harvested from de-novo auditors’ working dirs, deduplicated and lightly attributed. Orthogonal framing / pedagogy; off-ramp (scalar-vs-vector bound confusion) at the end. **Coverage:** one substrate reached a dedicated digested reflection (Gemini, AUDIT-WORKING-193847); a batched discrete/persistence walk (Codex, AUDIT-WORKING-613842) read it as part of the persistence cluster and judged it “doing serious work” with good scope honesty (its certified concern there is about the *downstream template/result* segments compressing the Model-S region-awareness, not this appendix).

1. Candidate Brief prose / pre-prose.

- **The discretization penalty as “measure twice, cut once.”** In a continuous world you can pull the hammer back mid-swing as the error crosses zero; in a discrete event-driven world the action commits the agent to the full consequence of the swing before the next observation arrives — so an over-high gain over-corrects and oscillates. A plain-language anchor for the no-overshoot bound $\eta^* < 2/c_{\max}$ (Gemini, AUDIT-WORKING-193847).

2. Candidate Discussion.

- **Latency bounds confidence, independent of epistemic certainty.** The no-overshoot condition means the *time between action and observation* caps the maximum safe update gain: a high-latency infrastructure forces the agent to lower η^* even when its model is perfectly certain ($U_M \approx 0$). “Confidence is not just a function of epistemic certainty; it is strictly bounded by the temporal resolution of the feedback loop.” A consequential framing of the discrete bound for agent/infrastructure design (Gemini, AUDIT-WORKING-193847).
- **The $1/\nu$ discretization penalty as a cost of event-driven agency.** The exact asymptotic gap $V_{ss} - V_\epsilon \approx n\sigma_w^2 c_{\max}^2 / (4c_{\min}^2 \nu)$ shows the discrete system is *systematically worse* than its continuous approximation by a factor scaling as $1/\nu$ under stochastic disturbance — high-frequency agents pay a smaller penalty. Names a real cost the continuous “fluid limit” hides; candidate to surface as the segment’s distinctive practical result (Gemini, AUDIT-WORKING-193847).

3. Follow-up items.

- **Heavy-tailed $\Rightarrow$ structural adaptation.** If w_k has infinite variance (Cauchy-style “black swan”), the mean-square Lyapunov proof fails entirely — implying such environments cannot be survived by parametric adaptation (η^* tuning) and *must* be met by structural adaptation (expanding R). Sharpens the existing heavy-tailed Working Note above into a routing claim toward #result-structural-adaptation-necessity (Gemini, AUDIT-WORKING-193847).

4. Readers often ask / wonder.

- **“Aren’t Lipschitz-bounded gradient descent and sector-bounded control theory different things?”** The segment’s quiet unification answer: they are the same physical reality at different temporal resolutions — gradient descent needs the stronger discrete Lipschitz bound (DA2’b) to prevent overshoot; continuous control needs only the directional sector bound (DA2’a) — surfacing this explicitly would help a reader coming from either side (Gemini, AUDIT-WORKING-193847).

Off-ramp (NOT gold) — routed for adjudication.

- **Scalar-vs-vector contraction-bound presentation is correct but confusing.** Gemini (AUDIT-WORKING-193847): the general (vector) stability condition $\eta^* < 2c_{\min}/c_{\max}^2$ and the scalar/colinear simplification $\eta^* < 2/c_{\max}$ disagree when $c_{\min} \ll c_{\max}$, and a reader can take this as a contradiction. The reconciliation: the general bound combines worst-case minimum alignment ($c_{\min}$) with worst-case maximum magnitude ($c_{\max}$) taken *independently* — necessary in high dimensions (orthogonal overshoot is possible) but overly pessimistic in 1D, where Cauchy-Schwarz is tight and all error is colinear overshoot. Suggested: add a sentence stating the general λ_{eff}^2 bound is conservative and that the gap between the two forms is the dimension artifact, not an inconsistency. Presentation/scope-honesty, not a math error.

Appendix P

Detail: Linear ODE Approximation

The full linear treatment of mismatch dynamics: scalar and vector forms, steady-state solutions under both disturbance models, transient behavior, convergence rate, validity conditions, breakdown modes, discrete-time connection, and adversarial coupling. This appendix collects the linear-case results that 3.9 states as a hypothesis, derives them explicitly, and delineates exactly where the linear approximation is valid and where it fails.

The detail serves as the *worked example* underlying the framework’s hypothesis that mismatch dynamics follow a first-order linear ODE in the simplest case. Everything here is exactly derivable from the linear assumption — the appendix’s purpose is to show, in detail, that the linear case gives clean closed-form results (the distinctive Model D $1/\mathcal{T}$ vs Model S $1/\sqrt{\mathcal{T}}$ scaling split; exponential transient convergence at rate $\mathcal{T}$; the squared-tempo law for adversarial coupling) and to *delineate exactly where the linear approximation is valid and where it fails*. The validity range: bounded mismatch within the sector region; non-saturating correction; sufficiently fast event rate relative to the mismatch timescale. The breakdown modes: saturation at large mismatch (correction slower than linear); threshold effects (small mismatches uncorrected — a dead zone); structural breakdown beyond a critical mismatch where the correction collapses entirely because the model class is no longer appropriate. The chapter’s central inequality — that adversarial mismatch ratios scale as the squared tempo ratio — is first introduced here under linearity, with the appendix showing exactly which assumptions the simple form rests on and where it transitions into the sector-condition machinery (4.4) as those assumptions relax.

Formal Expression.

1. The scalar and vector forms. The mismatch vector $\delta(t) \in \mathbb{R}^n$ evolves under the general dynamics from A:

Definition (vector dynamics).

$$\frac{d\delta}{dt} = -F(\mathcal{T}, \delta) + w(t) \quad (\text{P.1})$$

The **linear approximation** sets $F(\mathcal{T}, \delta) = \mathcal{T} \cdot \delta$, giving:

Hypothesis (linear-vector-ode).

$$\frac{d\delta}{dt} = -\mathcal{T} \cdot \delta + w(t) \quad (\text{P.2})$$

This is exact when the correction function is linear in δ (Kalman filter, Beta-Bernoulli conjugate update). It is an approximation otherwise.

Scalar form (equality). For a scalar mismatch $\delta(t) \in \mathbb{R}$:

Derived (scalar-ode, from linear-vector-ode with $n = 1$).

$$\frac{d\delta}{dt} = -\mathcal{T} \cdot \delta + w(t) \quad (\text{P.3})$$

This is an equality: the scalar ODE is exactly the $n = 1$ case of the vector ODE.

Norm form (inequality). For vector $\delta \in \mathbb{R}^n$, take the time derivative of $\|\delta\| = (\delta^T \delta)^{1/2}$:

$$\frac{d\|\delta\|}{dt} = \frac{\delta^T \dot{\delta}}{\|\delta\|} = \frac{\delta^T (-\mathcal{T} \delta + w)}{\|\delta\|} = -\mathcal{T} \|\delta\| + \frac{\delta^T w}{\|\delta\|} \quad (\text{P.4})$$

By Cauchy-Schwarz, $\delta^T w / \|\delta\| \leq \|w\|$. Under Model D (GA-2: $\|w(t)\| \leq \rho$):

Derived (norm-inequality, from Cauchy-Schwarz).

$$\frac{d\|\delta\|}{dt} \leq -\mathcal{T} \cdot \|\delta\| + \rho \quad (\text{P.5})$$

This is an inequality, not an equality. The Cauchy-Schwarz step is tight only when $w(t)$ is parallel to $\delta(t)$ (worst-case disturbance). The norm form stated in 3.9 is this upper bound. For the scalar case ($n = 1$), Cauchy-Schwarz is automatically tight and the inequality becomes an equality.

2. Steady-state solutions.

Model D (deterministic bounded disturbance, GA-2: $\|w(t)\| \leq \rho$). Setting $d\|\delta\|/dt = 0$ in the norm inequality gives the worst-case steady state:

Derived (model-d-steady-state, from norm inequality).

$$\|\delta\|_{ss} = \frac{\rho}{\mathcal{T}} \quad (\text{P.6})$$

This is the tight upper bound: steady-state mismatch equals the ratio of disturbance to correction. In the scalar case ($n = 1$), this is exact (the Ornstein-Uhlenbeck process with deterministic forcing converges to $\rho/\mathcal{T}$). In the vector case, it is the worst case over all disturbance directions.

Model S (stochastic zero-mean disturbance, GA-2S: $d\delta = -\mathcal{T} \delta dt + \sigma_w dW_t$). The Ornstein-Uhlenbeck process has stationary variance derived via Ito-Lyapunov analysis (Prop A.1S in A, specialized to $\alpha = \mathcal{T}$):

Derived (model-s-steady-state, from Ito-Lyapunov with linear correction).

$$\mathbb{E}[\|\delta\|^2]_{ss} = \frac{n\sigma_w^2}{2\mathcal{T}} \quad (\text{P.7})$$

The RMS steady-state mismatch is:

$$\|\delta\|_{rms} = \sigma_w \sqrt{\frac{n}{2\mathcal{T}}} \quad (\text{P.8})$$

For the scalar case ($n = 1$): $\|\delta\|_{\text{rms}} = \sigma_w/\sqrt{2\mathcal{T}}$.

Scaling difference. Model D gives $\|\delta\|_{ss} \propto 1/\mathcal{T}$; Model S gives $\|\delta\|_{\text{rms}} \propto 1/\sqrt{\mathcal{T}}$. Doubling the correction rate halves the deterministic steady-state mismatch but only reduces the stochastic steady-state by a factor of $\sqrt{2} \approx 1.41$. Correction is less effective against noise than against drift.

3. Transient solution. Model D, constant ρ . The linear ODE $d\|\delta\|/dt = -\mathcal{T}\|\delta\| + \rho$ with initial condition $\|\delta(0)\| = \|\delta_0\|$ has the standard first-order linear solution:

Derived (transient-model-d).

$$\|\delta(t)\| = \|\delta_0\| e^{-\mathcal{T}t} + \frac{\rho}{\mathcal{T}}(1 - e^{-\mathcal{T}t}) \quad (\text{P.9})$$

Mismatch decays exponentially from initial conditions toward the steady state $\rho/\mathcal{T}$ with time constant $\tau = 1/\mathcal{T}$. After k time constants, the transient has decayed by a factor of e^{-k} : 63% after one time constant, 95% after three, 99.3% after five.

Model S, Ornstein-Uhlenbeck. The mean-square mismatch evolves as (from Prop A.1S with $\alpha = \mathcal{T}$):

Derived (transient-model-s).

$$\mathbb{E}[\|\delta(t)\|^2] = \|\delta_0\|^2 e^{-2\mathcal{T}t} + \frac{n\sigma_w^2}{2\mathcal{T}}(1 - e^{-2\mathcal{T}t}) \quad (\text{P.10})$$

The variance converges twice as fast as the mean (rate $2\mathcal{T}$ vs. $\mathcal{T}$) because the Lyapunov function $V = \frac{1}{2}\|\delta\|^2$ is quadratic — its dynamics have doubled exponential rate.

4. When the linear approximation is valid. The linear case $F(\mathcal{T}, \delta) = \mathcal{T} \cdot \delta$ is a special case of the sector condition (GA-3) from 4.4 with:

Derived (linear-sector-parameters).

$$\alpha = \mathcal{T}, \quad R \rightarrow \infty \quad (\text{P.11})$$

That is, the sector condition holds globally with lower bound α equal to the tempo, and the sector-condition region has infinite radius. In the notation of 4.3:

- **Directional fidelity parameter:** $c_{\min} = 1$ (the correction points exactly at the mismatch).
- **Sector parameter:** $\alpha = \eta^* \cdot c_{\min} = \eta^* = \mathcal{T}/\nu$.
- **No upper saturation:** $c_{\max} = c_{\min} = 1$ (the sector bounds coincide).

Consequence for persistence. With $R \rightarrow \infty$, structural persistence ($\alpha > \rho/R$) is trivially satisfied for any $\mathcal{T} > 0$. The binding constraint is task adequacy alone: $\mathcal{T} > \rho/\|\delta_{\text{critical}}\|$ (Model D) or $\mathcal{T} > n\sigma_w^2/(2\|\delta_{\text{critical}}\|^2)$ (Model S). This is why the linear operational forms in 4.5 exist: in the linear world, structural persistence is free.

The linear approximation is exact when:

1. **Kalman filter** (scalar or matrix): the gain K_t produces $F(\delta) = K_t H \delta$, which is linear. The sector parameter equals the Kalman gain: $\alpha = K$ (scalar) or $\alpha = \lambda_{\min}^+(KH)$ (matrix). See 4.3.
2. **Beta-Bernoulli conjugate update**: the correction $F(\delta) = \delta/(n+1)$ is linear with $\alpha = 1/(n+1) = \eta_{\text{edge}}$.
3. **Exponential family with natural parameters**: the natural-parameter update is linear in the sufficient statistic, giving $\alpha = \eta \cdot \lambda_{\min}(\text{Fisher})$.
4. **Gradient descent on quadratic loss**: $F(\delta) = \eta \nabla^2 L \cdot \delta$ is linear with $\alpha = \eta \cdot \lambda_{\min}(\nabla^2 L)$.

5. When the linear approximation breaks down. The linear form $F = \mathcal{T} \cdot \delta$ fails when the true correction function deviates from linearity. Three failure modes, each corresponding to a specific nonlinearity:

Saturation at large $\|\delta\|$. The correction mechanism is overwhelmed by large mismatch. The true correction satisfies $F(\delta) < \mathcal{T}\|\delta\|$ for large $\|\delta\|$ — correction is slower than the linear prediction. Simulation confirms: for a saturating function $g(\delta) = \delta/(1 + |\delta|/R)$, the sector parameter at the capacity boundary is $\alpha \approx \mathcal{T}/2$ (half the linear value). The linear approximation overstates persistence margins. The sector-condition framework (4.4) handles this via finite R and reduced α .

Threshold effects at small $\|\delta\|$. Below a detection threshold ε , small mismatches go uncorrected: $F(\delta) \approx 0$ for $\|\delta\| < \varepsilon$. This creates a dead zone where the model drifts. The linear approximation (which predicts correction at all scales) misses this. The sector condition fails locally at small $\|\delta\|$ (the ratio $\delta^T F / \|\delta\|^2 \rightarrow 0$), but the dead zone is bounded and does not affect large-scale stability.

Structural breakdown. Beyond some critical mismatch, the correction rate drops to zero because the model class is no longer appropriate: the mismatch exceeds the model's representational capacity. $F(\delta) \approx 0$ for $\|\delta\| > R$. This is the structural adaptation trigger (4.8). The linear approximation, which predicts correction growing without bound, misses this entirely. The sector-condition framework captures it via the finite radius R of the sector-condition region.

6. Discrete-time connection. The continuous ODE is a fluid-limit approximation of the discrete event-driven dynamics (3.1). The discrete mismatch dynamics are:

$$\delta_{k+1} = \delta_k - \eta^* F_d(\delta_k) + w_k \quad (\text{P.12})$$

In the linear case ($F_d(\delta) = \delta$), this becomes:

$$\delta_{k+1} = (1 - \eta^*)\delta_k + w_k \quad (\text{P.13})$$

which is an AR(1) process with coefficient $\lambda = 1 - \eta^*$.

○ proves the formal connection:

Model D. The discrete and continuous steady states are *identical* — the fluid-limit gap is exactly zero. Both give $R^* = \rho/\alpha$. The discrete form has $\alpha = (1 - |\lambda|)/\eta^* = (1 - |1 - \eta^*|)/\eta^*$, which equals 1 (hence $\alpha_{\text{continuous}} = \nu \cdot 1 = \mathcal{T}$) when $0 < \eta^* < 1$.

Model S. The discrete steady-state variance is $\sigma_{\text{step}}^2/(1 - \lambda_{\text{eff}}^2)$, which recovers the continuous result $n\sigma_w^2/(2\mathcal{T})$ in the fluid limit ($\eta^* \rightarrow 0, \nu \rightarrow \infty, \nu\eta^* = \mathcal{T}$ fixed). The variance gap is $O(\eta^* c_{\text{max}}^2/c_{\text{min}}^2) = O(c_{\text{max}}^2/(c_{\text{min}}^2 \nu))$ — quantitatively small whenever $c_{\text{max}}^2/(c_{\text{min}}^2 \nu) \ll 1$. See ○ for the leading-order Taylor expansion that gives this scaling.

Fluid-limit validity. The ODE approximation is formally justified by the fluid limit theorem in ○ (conditional on Lipschitz regularity of F_d). The approximation is most accurate when $\eta^* \ll 1$ (the small-gain regime). It is least

accurate during initial transients when η^* is large, but this phase is short-lived and the transient error is bounded by $O(\eta^* c_{\max}/\nu^{1/2})$.

7. Adversarial coupling in the linear case. When two agents are coupled, A 's praxis contributes to B 's disturbance:

Hypothesis (Linear Coupling Model).

$$\rho_B = \rho_{B,\text{base}} + \gamma_A \cdot \mathcal{T}_A \quad (\text{P.14})$$

Under the linear approximation, the steady-state mismatches (Model D, coupling-dominant: $\gamma_A \mathcal{T}_A \gg \rho_{B,\text{base}}$) are:

Derived (linear-adversarial-steady-state, from Model D steady state + coupling model).

$$\|\delta_B\|_{ss} \approx \frac{\gamma_A \mathcal{T}_A}{\mathcal{T}_B}, \quad \|\delta_A\|_{ss} \approx \frac{\gamma_B \mathcal{T}_B}{\mathcal{T}_A} \quad (\text{P.15})$$

The ratio:

Derived (squared-tempo-law, from linear steady states).

$$\frac{\|\delta_B\|_{ss}}{\|\delta_A\|_{ss}} = \frac{\gamma_A}{\gamma_B} \left(\frac{\mathcal{T}_A}{\mathcal{T}_B} \right)^2 \quad (\text{P.16})$$

Under symmetric coupling ($\gamma_A \approx \gamma_B$), tempo advantage squares: a 2:1 tempo ratio yields a 4:1 mismatch ratio; a 3:1 ratio yields 9:1.

Why the squared law. The exponent 2 arises because the $1/\mathcal{T}$ steady-state scaling appears in both numerator (faster agent generates more disturbance) and denominator (faster agent corrects better): one power from the coupling, one from the steady-state. Under Model S (stochastic coupling, $1/\sqrt{\mathcal{T}}$ scaling), the exponent reduces to 3/2. See 16.3 for the full regime analysis.

Simulation validation. Variant A confirmed $b = 1.999$ (Model D, coupling-dominant). Variants C-D confirmed $b = 1.481$ (Model S, coupling-dominant). Both within numerical precision of the derived values. See AX.

Epistemic Status.

Exact for the linear case. The linear ODE is a standard first-order linear system; the solutions (transient and steady-state, both Model D and Model S) are textbook results for the Ornstein-Uhlenbeck process and its deterministic analog. The sector-condition parameters ($\alpha = \mathcal{T}, R \rightarrow \infty$) are exact when the correction function is linear. The discrete-time connection is exact (zero gap for Model D, $O(\eta^*)$ gap for Model S — see O for the leading-order Taylor expansion). The adversarial scaling law ($b = 2$ for Model D, $b = 3/2$ for Model S) is derived and simulation-validated.

What is exact vs. what is an approximation. The linear ODE is exact for systems with linear correction dynamics (Kalman, Beta-Bernoulli, quadratic loss). For nonlinear systems, the linear ODE is a local Taylor approximation valid near the steady state. The sector-condition framework (4.4, A) is the general treatment; this appendix derives what falls out when the sector bounds coincide.

The norm form is an upper bound for $n > 1$. The inequality introduced by Cauchy-Schwarz at the norm step is tight only in the worst case. For typical disturbances that are not persistently aligned with the mismatch direction, the actual norm trajectory lies below the upper bound. The scalar case ($n = 1$) is an equality. This distinction is load-bearing:

downstream results that use $d\|\delta\|/dt = -\mathcal{T}\|\delta\| + \rho$ as an equality (e.g., the exact transient solution, the exact steady-state ratio) are strictly correct only for $n = 1$ or for worst-case disturbance. The Lyapunov analysis in A avoids this issue by working with $V = \frac{1}{2}\|\delta\|^2$ directly, without passing through the norm derivative.

Max attainable: exact (the linear case is its own exact analysis; there is nothing to tighten).

Discussion.

Why this appendix exists. The linear ODE in 3.9 is labeled as a hypothesis and described as a “first-order approximation.” That is correct for the general case. But for agents with linear correction dynamics, the ODE is not an approximation at all – it is exact. This appendix serves two roles: (1) providing the complete derivation that 3.9 points to rather than carrying inline, and (2) making precise when the “approximation” is exact and when it genuinely introduces error.

Pedagogical value. The linear case is the entry point for understanding AAT’s mismatch dynamics. The closed-form solutions (exponential convergence, ratio steady state, squared scaling law) build intuition that the general sector-condition results (4.4) cannot provide – Lyapunov gives bounds and existence, not quantitative trajectories. The linear treatment also connects directly to the AR(1) simulation framework (AX) and classical control theory (PI controllers, Kalman filters).

Speed-quality substitutability. From $\mathcal{T} = \nu \cdot \eta^*$ (single channel): doubling event rate ν has the same effect on $\|\delta\|_{ss}$ as doubling update quality η^* . They are multiplicative when both improve: 50% improvement in each yields $1.5 \times 1.5 = 2.25\times$, not $3\times$. This is the formal analog of Boyd’s insight that Orient quality often matters more than raw OODA speed.

The persistence threshold in the linear world. From the steady state: $\|\delta\|_{ss} < \|\delta\|_{critical}$ iff $\mathcal{T} > \rho/\|\delta\|_{critical}$ (Model D). This is task adequacy only – structural persistence is trivially satisfied because $R \rightarrow \infty$. The linear world makes persistence look easier than it is. For nonlinear agents, both structural and task-adequacy conditions must hold (4.5).

Working Notes.

- The vector ODE $d\delta/dt = -\mathcal{T}\delta + w$ assumes isotropic correction (same rate in all directions). For anisotropic agents (tempo tensor $\mathbf{T}$), the linear ODE generalizes to $d\delta/dt = -\mathbf{T}\delta + w$ with matrix exponential solutions and per-eigenvalue convergence rates. The per-dimension persistence condition ($\mathcal{T}_k > \rho_k/\delta_{critical,k}$) from 16.5 is the diagonal special case.
- The norm-inequality subtlety (Section 1 above) propagates into the adversarial scaling law: the squared law ($b = 2$) is exact for $n = 1$ and a worst-case upper bound for $n > 1$. Whether this matters quantitatively for multi-dimensional adversarial dynamics has not been tested in simulation.
- Model S’s Ito-Lyapunov derivation assumes σ_w is constant. When the noise scale depends on state ($\sigma_w(\delta)$, as in multiplicative noise models), the Ornstein-Uhlenbeck analysis no longer applies and the steady-state formula changes. The sector-condition results in A handle additive noise only.

Incidental audit gold (gold-lift, 2026-05-31). Cross-audit ideation harvested from a de-novo auditor’s working dir, lightly attributed. Orthogonal framing / pedagogy; off-ramp (steady-state tightness) at the end. **Coverage:** one substrate reached a digested reflection (Gemini, AUDIT-WORKING-193847), which read the segment as both pedagogical and rigorous and praised the failure-mode mapping.

1. Candidate Briefprose / pre-prose.

- **The “Linearity Illusion.”** The linear approximation assumes $R \rightarrow \infty$ and therefore *cannot see* the structural-breakdown trigger — “linear math makes survival look easier than it is.” A compact framing of why the linear ODE is the right *baseline* but the wrong *safety analysis*: it gives closed-form trajectories in friendly (convex / linear) environments against which the cost of nonlinearity can be measured (Gemini, AUDIT-WORKING-193847).

2. Candidate Discussion.

- **Dead zones as a *requirement for sanity*, not a defect.** The threshold effect at small $\|\delta\|$ (errors below sensor resolution go uncorrected, so the agent drifts) means absolute perfection is impossible for any finite-resolution agent — but more deeply, *correcting* a tiny error often costs more coordination overhead C than the error itself costs, so a rational agent *builds* a dead zone. For an ELI: a zero-tolerance infrastructure would be paralyzing; “a mind that constantly updates its worldview based on microscopic sensor noise will consume all its tempo on pointless neuroses.” Tolerance for minor mismatch as a design principle — vivid Discussion candidate (Gemini, AUDIT-WORKING-193847).
- **Speed-quality substitutability and the 2.25× tempo multiplier.** Since $\mathcal{T} = \nu \cdot \eta^*$, a 50% gain in speed *and* a 50% gain in quality compound to a 2.25× tempo improvement — a clean illustration of why tempo is multiplicative in its two factors, with the (open) investment-strategy implication that balancing compute (ν) against architectural quality (η^*) may beat pouring everything into one (Gemini, AUDIT-WORKING-193847).
- **Nonlinear-failure-mode mapping.** The “when the linear approximation breaks down” section (dead zones, saturation, structural breakdown) was singled out as one of the clearest expositions of nonlinear control failure the auditor had read, with each mode mapped to a specific AI failure — candidate to foreground as the segment’s pedagogical payload (Gemini, AUDIT-WORKING-193847).

3. Follow-up items.

- The existing anisotropic-correction Working Note (matrix tempo $\mathbf{T}$) was implicitly re-affirmed; no new follow-up beyond what is already noted.

4. Readers often ask / wonder.

- **“Is $\rho/\mathcal{T}$ a loose worst-case bound or the real steady state?”** See off-ramp — the answer (for constant drift it is exactly tight at $t \rightarrow \infty$) is worth stating so a reader does not under-weight it as a rarely-hit bound (Gemini, AUDIT-WORKING-193847).

Off-ramp (NOT gold) — routed for adjudication.

- **The Model-D steady state $\rho/\mathcal{T}$ is exactly tight (not merely a worst-case bound) for constant directional drift.** Gemini (AUDIT-WORKING-193847): the norm-form $d\|\delta\|/dt \leq -\mathcal{T}\|\delta\| + \rho$ is an inequality because Cauchy-Schwarz is tight only when the disturbance w aligns with δ — but for a *constant* drift vector, a linear system’s $\delta(t)$ eventually aligns with w , making the inequality an equality, so $\|\delta\|_{ss} = \rho/\mathcal{T}$ is the exact inevitable steady-state magnitude for any persistent directional drift (Model D, $n > 1$), not a rarely-hit bound. Suggested: clarify that the transient is strictly bounded by the inequality while the steady state is practically tight at $t \rightarrow \infty$. Scope/tightness clarification — strengthens the result’s certainty rather than weakening it.

Appendix Q

Derivation: Graph Structure Uniqueness

Operational requirements on the agent’s representation — directed temporal ordering, probabilistic uncertainty, and the ability to test strategy components — are *sufficient* for the strategy to be a directed acyclic graph with the Markov factorization property. The argument parallels Cox’s theorem for probability in *form* but not yet in *strength*: Cox’s theorem is necessary-and-sufficient (the only measure satisfying the desiderata is probability); this result is sufficient only (the desiderata guarantee DAG+Markov, but no one has shown a non-DAG structure cannot satisfy them). Acyclicity and directed edges are *proved* from temporal ordering over a finite horizon; the Markov factorization is *proved under causal sufficiency* via the Causal Markov Condition theorem (Spirtes–Glymour–Scheines, Pearl). A fourth postulate (observable intermediates) is required for localized strategic diagnosis but not for the representation or persistence results themselves.

How far the Cox parallel goes. Cox’s theorem starts from desiderata on how a rational agent should quantify uncertainty (consistency, universality, continuous functional composition) and proves that the *only* measure satisfying them is probability — necessity. This segment starts from desiderata on how a bounded agent can represent its strategy under causal action (directed temporal order, probabilistic edge uncertainty, causal sufficiency of the chosen nodes) and proves that DAG+Markov *suffices* to satisfy them — sufficiency. The necessity direction — no non-DAG structure (factor graphs, junction trees with cyclic message schedules, chain graphs) can satisfy P1–P4 plus causal sufficiency — is not established here. In practical terms this gap is unimportant because the proved sufficiency gives a rigorous grounding for the DAG structure; claims that AAT’s strategy *must* be a DAG should be read as “must-if-sufficient-via-this-route,” not as a proved necessity. A stronger Cox-style result is open. The placement of the DAG structure on a footing comparable to probability — a consequence of operational requirements rather than a modeling convenience — holds in the sufficiency direction; the full parallel with Cox awaits a uniqueness argument.

Formal Expression.

The Postulates. Four properties that a strategy representation must satisfy. Each is independently motivated from the adaptive-systems foundation.

P1: Directed Temporal Ordering.

If component *A* of the strategy causally produces component *B*, then *A* temporally precedes *B*. The strategy representation must respect this directionality — edges point from causes to effects, from actions to outcomes, from prerequisites to goals.

Derived (from #post-causal-structure).

This is a consequence of the temporal postulate (1.7): the arrow of time is constitutive, not incidental. Reversing a causal edge would mean effects precede causes, which is physically impossible.

P2: Probabilistic Uncertainty.

The agent's uncertainty about whether each step of the strategy will succeed must be quantified by a measure satisfying Cox's axioms (consistency, universality, non-negativity). The unique such measure is probability (Cox 1946, "Probability, Frequency and Reasonable Expectation," *American Journal of Physics* 14(1):1–13; modern exposition in Jaynes 2003, *Probability Theory: The Logic of Science*, Cambridge University Press, Chapter 2). The agent may use other representations internally (confidence scores, fuzzy logic), but these must be mappable to probability to be consistent.

Derived (from Cox's theorem).

P3: State-Local Revisability.

When the agent observes evidence about one component of its strategy (e.g., "step 3 succeeded" or "prerequisite 2 is blocked"), it must be able to update its beliefs about that component and its consequences without recomputing the entire strategy from scratch.

Derived (from #der-chain-confidence-decay + bounded computation).

Why this is forced, not chosen:

From fragility. Additive log-confidence (8.1) means longer chains are exponentially less reliable. The agent will frequently encounter partial failures. Each partial failure requires re-evaluation of the affected portion of the strategy. If each re-evaluation requires full recomputation, the agent's planning tempo T_Σ is catastrophically slow — potentially violating the strategy persistence condition.

From bounded computation. The agent has finite computational resources (the IB constraint applies to planning as well as model maintenance). Full recomputation of a strategy with N components costs $O(N)$ or worse. Local revision costs $O(|\text{affected}|)$, which can be much smaller.

From the persistence condition. Strategy must be revised faster than the environment invalidates it. Local revision directly increases T_Σ by reducing the per-update cost. An agent that must recompute everything on each update has lower T_Σ and is more likely to fall below the persistence threshold.

P4: Observable Intermediates.

To support **localized strategic diagnosis and revision**, the strategy representation benefits from internal checkpoints — observable states between the initial action and the final goal — that the agent can monitor to detect partial failure.

Derived (from #der-chain-confidence-decay + monitoring requirement).

Without intermediates, the agent cannot detect chain failure until the final outcome. By the time the final outcome reveals failure, all intermediate actions are wasted. With intermediates, the agent can detect failure at step k and revise, saving the cost of steps $k + 1$ through n . The value of early detection grows with chain length, because longer chains fail more often (P2 + 8.1).

Observable intermediates are not required for strategy representation or persistence. When intermediates are unobservable, plan-level tracking (9.10, Case 3) preserves the sector condition at the cost of per-edge diagnostic resolution — the agent knows the plan is failing but cannot localize which step needs revision (9.3). P4 is therefore a requirement for *strong diagnostics*, not for strategy representation per se. The observability investment tradeoff (9.3) quantifies the payoff: making an intermediate observable improves the sector parameter from $1/(n_\Phi + 1)$ (plan-level) to $\min(1/(n_1 + 1), \theta_1/(n_2 + 1))$ (per-edge weakest-link).

The Derivation.

Step 1: P1 implies directed edges. Each component X_i of the strategy has a set of direct causes $\text{Pa}(X_i)$ — the components whose outcomes directly influence X_i 's outcome. P1 requires that these causal relationships are directed: $\text{Pa}(X_i)$ temporally precedes X_i . This gives directed edges $\text{Pa}(X_i) \rightarrow X_i$.

Step 2: P2 implies probability distributions on edges. Each edge carries uncertainty: $P(X_i | \text{Pa}(X_i))$. By P2 (Cox), this is a probability distribution. The joint distribution over all strategy components is some $P(X_1, \dots, X_n)$.

Step 3: Causal sufficiency implies the Markov condition (proved).

Claim. For a causally sufficient strategy DAG, the Markov factorization property is a theorem — a consequence of the Causal Markov Condition (CMC).

Derived (Conditional on causal sufficiency of Σ_t).

The Markov factorization property. Each variable X_i is conditionally independent of its non-descendants given its parents:

$$X_i \perp \text{NonDesc}(X_i) \mid \text{Pa}(X_i) \quad (\text{Q.1})$$

Equivalently, the joint distribution factorizes as:

$$P(X_1, \dots, X_n) = \prod_{i=1}^n P(X_i \mid \text{Pa}(X_i)) \quad (\text{Q.2})$$

(The equivalence holds for positive distributions — Lauritzen 1996, Theorem 3.27.)

The argument has five parts:

(a) The DAG is a causal model. P1 establishes that edges represent causal relationships: completing a parent step causally advances the child step. P2 establishes probabilistic uncertainty over outcomes. Together: Σ_t is a causal DAG in the sense of structural causal models (Pearl 2009, Definition 7.1.1) — each node's outcome is determined by its parents' outcomes (through the causal mechanism encoded in the edge credences) plus exogenous uncertainty specific to that step. Formally, each node admits a structural equation:

$$X_i = f_i(\text{Pa}(X_i), \varepsilon_i) \quad (\text{Q.3})$$

where f_i is the local causal mechanism and ε_i is the exogenous noise (the residual uncertainty at step i not determined by its parents).

(b) Causal sufficiency implies exogenous independence. The exogenous terms ε_i are mutually independent if and only if no unmodeled common cause affects two or more nodes in the graph. This is precisely the **causal sufficiency** assumption: every variable that is a direct common cause of two or more nodes in Σ_t is itself a node in Σ_t .

For agent-constructed strategies, causal sufficiency is a **modeling ideal, not a typical condition**. The agent designed the graph, so all *intended* causal relationships are explicit — but environmental common causes (shared infrastructure, weather, market shifts, correlated adversary actions) routinely affect multiple strategy steps without appearing as nodes. In complex, multi-stakeholder, or adversarial environments, causal insufficiency is the dominant case (8.3, Correlation Hierarchy). When an environmental factor is omitted, the exogenous terms become correlated and the Markov condition fails. This is model inadequacy (4.8), and the remedy is to add the missing common-cause node — but identifying which common causes matter is a modeling judgment, not a mechanical procedure (8.3, L1 construction principle). The

proof's conditional on causal sufficiency is therefore a condition on model quality: the result holds exactly when the DAG is well-constructed, approximately when it is close, and fails when major common causes are missing. The Correlation Hierarchy in 8.3 provides the practical framework: L0 (independence, this proof's assumption) gives tractable results; L1 (augmented DAG with explicit common-cause nodes) is the practical default in complex domains; L0 formal results transfer to correctly constructed L1 DAGs.

(c) **The Causal Markov Condition theorem.** For a DAG G over variables $V = \{X_1, \dots, X_n\}$ with structural equations $X_i = f_i(\text{Pa}(X_i), \varepsilon_i)$ where the ε_i are mutually independent:

$$P(X_1, \dots, X_n) = \prod_{i=1}^n P(X_i \mid \text{Pa}(X_i)) \quad (\text{Q.4})$$

This is the **Causal Markov Condition** — a proved theorem, not a modeling assumption. The standard references are Spirtes, Glymour, and Scheines (2000, Theorem 3.4) and Pearl (2009, §1.4.1, Theorem 1.4.1). The proof applies the chain rule in topological order: $P(X_1, \dots, X_n) = \prod_i P(X_i \mid X_1, \dots, X_{i-1})$, then uses the independence of ε_i to show that conditioning on all predecessors reduces to conditioning on parents only. Each non-parent predecessor's influence on X_i is fully mediated through the parents — its direct contribution enters through the causal mechanism f_i , not through ε_i .

(d) **P3 as consequence.** State-local revisability (P3) was originally stated as an independent postulate. The CMC reveals it is a *consequence* of the causal structure under causal sufficiency: since $X_i \perp \text{NonDesc}(X_i) \mid \text{Pa}(X_i)$, updating beliefs about X_i requires only $\text{Pa}(X_i)$ — local revision is automatically correct. No information from the rest of the graph changes the conditional distribution of X_i given its parents. P3 was motivated as an operational requirement (agents *need* local revision for computational tractability, and the persistence condition demands it). The CMC shows the requirement is automatically satisfied by any causally sufficient causal DAG. The two arguments converge from different directions: P3 says local revision is *needed*; the CMC says it is *guaranteed* (under causal sufficiency).

(e) **Connection to edge independence.** The CMC's exogenous independence condition (ε_i mutually independent) is precisely the **edge-independence assumption** in the AND/OR status propagation (8.3). When exogenous noise terms are independent, edge outcomes are conditionally independent given parents, and the AND/OR formulas compute correct probabilities. When they are correlated (causal insufficiency — latent common causes), the AND/OR propagation systematically overestimates success because it treats joint failure probability as the product of marginals. The validity of the Markov factorization and the validity of the independence model are the *same condition*: causal sufficiency of Σ_t . See 8.3 for the full treatment of correlated failure as the primary case.

Assembling (a)-(e). P1-P2 establish that Σ_t is a causal DAG with probabilistic uncertainty. Under causal sufficiency (exogenous independence), the CMC theorem proves the Markov factorization. P3 (local revisability) follows as a validated consequence. The Markov property is both operationally required (P3) and structurally guaranteed (CMC). When causal sufficiency fails, the Markov factorization is still the agent's *intended* factorization — the one its DAG represents — but it is wrong about the world. The gap between intended and actual factorization manifests as correlated failure and $\hat{R}_2$ overestimation, and the fix is structural: add the missing common-cause nodes to restore causal sufficiency.

Step 4: P1 + finite horizon implies acyclicity (proved). This is the strongest piece of the argument. See the dedicated section below.

Step 5: Assembly. P1 (directed edges + causal interpretation) + P2 (probabilistic) + causal sufficiency (CMC $\rightarrow$ Markov factorization) + P4 (internal nodes) + finite horizon (acyclicity):

The strategy representation must be representable as a directed acyclic graph with probability distributions at each node conditioned on its parents — a Bayesian network. P3 (local revisability) is validated as a consequence of this structure, not required as a premise.

Acyclicity Derivation.

This resolves a former known fragility in the theory. Acyclicity of Σ_t is derived, not assumed.

Derived (from #post-causal-structure + finite planning horizon).

Result. For a strategy representation over a finite future horizon, temporal ordering forces acyclicity.

Derivation.

1. Each node X_i in Σ_t represents a future event or state with temporal position $\tau_i > t$ (the future time at which the step occurs or the state is evaluated).
2. Each edge $X_i \rightarrow X_j$ requires $\tau_i < \tau_j$ (P1: causes temporally precede effects).
3. A cycle $X_i \rightarrow X_j \rightarrow \dots \rightarrow X_i$ would require $\tau_i < \tau_j < \dots < \tau_i$, which is impossible for a real-valued time index.
4. Therefore the graph is acyclic. $\square$

Formally: a finite set with a strict partial order (future events ordered by time) is representable as a DAG. This is a standard result in order theory — every finite partial order has a Hasse diagram, which is a DAG.

The iteration objection resolved. A strategy that says “try A , if fail try B , if fail try A again” appears cyclic.

In the time-indexed representation:

$$A_1 \rightarrow \text{check}_1 \rightarrow B_1 \rightarrow \text{check}_2 \rightarrow A_2 \rightarrow \dots \quad (\text{Q.5})$$

Each attempt is a distinct node at a distinct time. The apparent cycle is a linear chain in the unrolled view. Iteration “terminates” when either a node succeeds (remaining retry nodes become probability-zero), the agent exhausts its resource budget (a constraint truncating the chain), or the horizon ends. Any finite-horizon strategy, including those with “loops” in the informal sense, is acyclic when time-indexed.

Scope. This applies to Σ_t (the agent’s strategy over the future), not to M_t ’s model of the environment. M_t may include cyclic causal processes — feedback loops in the physical world, market dynamics, ecosystem interactions. The acyclicity is specific to the purposeful substate because Σ_t represents planned future actions and the future is partially ordered by time. M_t ’s model of environmental dynamics may need to represent cycles (via time-unrolled DBNs or other cyclic structures).

Connection to Pearl. Pearl’s do-calculus is defined on DAGs. Extensions to cyclic structures exist (cyclic SCMs, equilibrium models) but are substantially more complex and lose some of do-calculus’s clean properties. The temporal argument here shows that for strategy representations (future-looking plans), acyclicity is not a convenience restriction on Pearl’s framework — it is a consequence of the temporal structure of planning.

What Is Derived vs. What Is Chosen.

Table Q.1

<i>Property</i>	<i>Motivating postulate</i>	<i>Strength</i>
Directed edges	Temporal ordering (P1, 1.7)	Proved
Probabilistic uncertainty	Cox’s theorem (P2)	Proved
Acyclicity	Temporal ordering + finite horizon (P1)	Proved
Internal structure	Fragility + monitoring (P4, 8.1)	Derived

<i>Property</i>	<i>Motivating postulate</i>	<i>Strength</i>
Markov factorization	Causal Markov Condition theorem (P1 causal interpretation + P2 probability + causal sufficiency)	Proved under causal sufficiency (CMC theorem)
DAG with Markov property	P1 + P2 + causal sufficiency (CMC) + P4	Conditional on causal sufficiency — which is testable and repairable
AND/OR parameterization	Boolean completeness + parsimony	Hypothesis (binary outcomes only)
Single-parameter edges	Parsimony / IB	Formulation choice
Specific node ontology	—	Formulation choice

The dividing line: acyclicity and directed edges are proved; the full DAG-with-Markov-property is conditional on causal sufficiency. The parameterization (AND/OR, CPT form, edge semantics) is a formulation choice within the strongly motivated structure, motivated by parsimony and domain fit but not by mathematical necessity.

Equivalence Class. Within the DAG class. Multiple DAGs can encode the same conditional independence relations, forming a Markov equivalence class identified by a CPDAG (completed partially directed acyclic graph). Two DAGs in the same equivalence class make identical probabilistic predictions but may differ in causal interpretation.

Across representation types. Factor graphs, junction trees, influence diagrams, and chain graphs are NOT simple presentational variants of DAGs:

- **Factor graphs** and **junction trees** preserve factorization and inference structure without necessarily preserving directed causal semantics.
- **Influence diagrams** add decision and utility nodes — a richer object, not an equivalent one.
- **Chain graphs** can express independence models that are not representable as DAGs at all.
- **Markov equivalence** is a statement within DAG classes, not across all graphical model types.

The correct claim is narrow: for a given factorized distribution, DAG and factor-graph representations can compute the same marginals. But causal semantics (do-calculus) are DAG-specific and do not transfer to undirected or mixed representations without additional structure.

AAT's choice. AAT uses DAG + AND/OR because: (a) AND/OR is the most parsimonious complete basis for binary combination (8.2), (b) the DAG naturally supports causal/interventional reasoning (Pearl's do-calculus), and (c) the representation converged across three independent formalism attempts.

Epistemic Status.

The acyclicity derivation is *exact* — it follows from temporal ordering over a finite horizon via standard order theory. The individual postulates P1, P2, and P4 are each well-grounded (temporal structure, Cox's theorem, and chain fragility respectively).

The Markov property is now *proved under causal sufficiency* via the Causal Markov Condition theorem (Spirtes, Glymour & Scheines 2000, Theorem 3.4; Pearl 2009, Theorem 1.4.1). The previous sketch argument (P3 requires locality → P1 identifies parents → therefore Markov) has been replaced by a rigorous chain: P1-P2 establish the causal DAG structure, causal sufficiency guarantees exogenous independence, and the CMC theorem proves the factorization. P3 (local revisability) is now a *consequence* of the Markov property, not a premise — the CMC shows that local revision is automatically correct for causally sufficient DAGs, validating P3's operational requirement.

The conditioning on causal sufficiency is the right level of conditionality: it is a property of *strategy quality* (did the agent include all relevant common causes as nodes?), not of agent architecture. It is testable in principle (correlated

residuals after convergence indicate missing common causes) and repairable in practice (add the common-cause node). When causal sufficiency fails, the Markov factorization fails, and this manifests as correlated failure and $\hat{P}_2$ overestimation (8.3). The edge-independence assumption in AND/OR propagation and the causal sufficiency condition for the Markov property are the same condition viewed from different angles.

Max attainable: *exact* for acyclicity (already there). *Derived conditional* for the full DAG-with-Markov-property — the derivation is rigorous (invokes a proved theorem), and the remaining condition (causal sufficiency) is about model quality, not proof quality.

The AND/OR restriction is a *hypothesis* for binary outcomes, grounded in Boolean completeness and parsimony. For non-binary outcomes, it does not apply and richer parameterizations within the derived graphical structure are needed.

The parallel to Cox’s theorem is now tighter than previously stated: Cox’s theorem proves that consistency axioms force probability; the CMC theorem proves that causal structure under sufficiency forces the Markov factorization. Both are formal results, not analogies. The remaining gap: Cox’s axioms are necessary and sufficient for probability; AAT’s postulates are sufficient for DAG+Markov structure, but the necessity direction (could a non-DAG structure satisfy P1-P4?) is not established. For practical purposes this gap is unimportant — the proved sufficiency gives a rigorous foundation for the strategy representation.

Discussion.

The contribution of this analysis. The dividing line between derived/conditional structure and chosen parameterization is the key result. The agent community (and future researchers) know exactly what is proved (acyclicity, directed edges), what is conditional (Markov property, contingent on causal sufficiency), and what is a formulation choice (the parameterization within the graph). Alternative parameterizations can be explored without abandoning the theoretical foundation.

Acyclicity deserves emphasis. The existing theory previously flagged “DAG acyclicity is an assumption, not forced” as a known fragility. The derivation above resolves this: acyclicity of Σ_t follows from temporal ordering over a finite planning horizon. This is specific to the strategy — M_t ’s environment model is not restricted to acyclic structures.

The causal sufficiency assumption. The Markov condition assumes causal sufficiency — no hidden common causes within the strategy. Since the agent constructs Σ_t , there are no “hidden” variables internal to the strategy (the agent chose all the nodes). But environmental common causes that affect multiple strategy steps may violate the Markov condition within Σ_t . If so, latent variable models or causal graphs with hidden nodes would be needed. This connects to the edge identifiability problem noted in 8.3.

Working Notes.

- **P3→Markov: resolved.** The previous gap (P3 could be satisfied by non-Markov sparse factorizations) is resolved by grounding the Markov property in the CMC theorem rather than in P3. The CMC proves the factorization from the causal structure (P1) and causal sufficiency (exogenous independence), making P3 a consequence rather than a premise. The alternative-factorization concern (factor graphs, junction trees) is now moot: the Markov property follows from the causal semantics, not from locality. P3 remains valuable as an *operational requirement* that the CMC-guaranteed factorization satisfies.
- **Parsimony theorem for AND/OR.** Is there a formal result that AND/OR (noisy-AND + noisy-OR) is the unique $O(k)$ -parameter complete basis for binary nodes? This would strengthen 8.2 from formulation choice to derived. The Boolean completeness half is standard; the uniqueness-under-parsimony half is not established.
- **Non-binary outcome analogs.** For continuous or multi-valued outcomes, the natural analogs of AND/OR might be min/max or additive/multiplicative combination. Is there a completeness result for these? The current argument applies cleanly only to binary outcomes.

- **Environmental common causes and the CMC.** The CMC proof makes the failure mode precise: when unmodeled environmental dependencies (weather, shared infrastructure, market conditions) affect multiple strategy steps, the exogenous noise terms ϵ_i become correlated, causal sufficiency fails, and the Markov factorization is violated. The consequence is exactly the correlated-failure phenomenon in 8.3: $\hat{R}_E$ overestimates success because it treats joint failure probability as the product of marginals. The fix is structural: add the common-cause as a condition node in Σ_t , restoring causal sufficiency and the Markov property for the augmented DAG. This connects the graph-theoretic result (Markov property) to the strategy-layer result (independence model validity) through a single condition: causal sufficiency.
- **Promotion potential.** With the CMC-based proof, 8.3's DAG structure claim is now strongly grounded: acyclicity is proved, the Markov property is proved under causal sufficiency. The remaining “definition” character of strategy-dag is about the *parameterization* (AND/OR, single-parameter edges), which is a formulation choice within the proved graphical structure.

Incidental audit gold (gold-lift, 2026-05-31). Cross-audit ideation harvested from de-novo auditors' working dirs, deduplicated and lightly attributed. Orthogonal framing / pedagogy; off-ramp (time-unrolled scope clarification) at the end. **Coverage:** two substrates reached a dedicated digested reflection (Gemini, AUDIT-WORKING-193847; Gemini-voiced, AUDIT-WORKING-773921); three batched strategy-DAG walks touched it by reference and converged on the same praise (Claude, AUDIT-WORKING-471203; Claude, AUDIT-WORKING-963715; Claude, AUDIT-WORKING-451729).

1. Candidate Briefprose / pre-prose.

- **“The Markov property in planning is a guaranteed consequence of a causally-sufficient world-model, not an arbitrary RL convenience.”** The recurring one-line synthesis: many planning frameworks (A, *MCTS*) *assume* DAG structure and Markov factorization; here acyclicity is derived from temporal ordering and the factorization from the Causal Markov Condition under causal sufficiency. A plain-language anchor (Gemini, AUDIT-WORKING-193847; Gemini-voiced, AUDIT-WORKING-773921).
- **“You can't cause your own prerequisite.”** Acyclicity in one breath: a strategy is about *causing future events*, and a finite-horizon plan is a finite poset ordered by time, representable as a DAG (its Hasse diagram) — the same argument as the acyclicity of Pearl's causal DAGs, applied to strategy. Named as “a quiet structural contribution that's easy to miss” (Claude, AUDIT-WORKING-963715; the poset/Hasse framing, AUDIT-WORKING-773921).

2. Candidate Discussion.

- **The Cox's-theorem-parallel as an epistemic-humility exemplar.** The segment proves the DAG structure is *sufficient* (a Markov DAG meets all operational requirements) but explicitly flags *necessity* (no other structure could) as open, parallel to “Cox's theorem in form but not yet in strength.” Two substrates read this self-restraint as a trust-building move worth preserving prominently — and the “What Is Derived vs. What Is Chosen” table was independently called “the platinum standard for epistemic clarity,” pre-empting any “you just assumed a DAG” critique by showing exactly which physics/information postulates force it (Gemini-voiced, AUDIT-WORKING-773921; Gemini, AUDIT-WORKING-193847).
- **Local revisability (P3) demoted from premise to derived consequence.** Auditors flagged the refactor that turns Postulate 3 into a *consequence* of the CMC as “stunning structural refactoring” / “the theory simplifying and tightening itself” — candidate to surface as a worked example of the framework tightening its own axiom set (Gemini, AUDIT-WORKING-193847; Gemini-voiced, AUDIT-WORKING-773921).
- **“Paranoia” as the active search for hidden common causes.** The causal-sufficiency failure mode (omitting environmental common causes correlates the exogenous ϵ_i , so $\hat{R}_E = \prod p_i$ becomes wildly optimistic) glossed as: an agent naturally drifts to L0 independence because it is cheaper, so it must be *forced* to compute the covariance of its failures $\hat{p}_{ij}$ to detect missing L1 common causes. “A mind without paranoia is mathematically guaranteed to be blindsided by correlated failure” — a vivid framing of the existing common-cause Working Note (Gemini, AUDIT-WORKING-193847).

3. Follow-up items.

- **AND/OR parsimony / uniqueness.** Both the segment's own Working Note and an auditor flag the same open: if AND/OR (noisy-AND + noisy-OR) can be proven the *unique* $O(k)$ -parameter complete basis for binary combination, #scope-and-or upgrades from formulation-choice to derived — “locking the entire strategy structure into pure math” (Gemini-voiced, AUDIT-WORKING-773921).
- **Necessity direction stays watched downstream.** Auditors note a consistency-watch: downstream segments should not treat the DAG as the *only* possible strategy representation while necessity is open (Claude, AUDIT-WORKING-451729).

4. Readers often ask / wonder.

- **“What forces the agent to fix a missing common cause?”** The CMC framing makes the trigger precise but uncomfortable: causal-sufficiency violations show up as *correlated residuals after convergence*, so the agent must experience repeated inexplicable joint failures before it can trigger L1 structural adaptation to add the missing node — a reader may want this failure-then-repair loop made explicit (Gemini, AUDIT-WORKING-193847).

Belongs elsewhere.

- **The L0/L1/L1'/L2 correlation hierarchy belongs to #def-strategy-dag, not this appendix.** Auditors locate the richer correlation-hierarchy content (direction/magnitude of bias; the “structurally refuted” L1'-with-unobservable-common-cause identifiability obstruction) in #def-strategy-dag; this appendix carries the acyclicity + CMC-factorization proofs. Noted so a future promoter routes correlation-hierarchy gold to the definition segment (Claude, AUDIT-WORKING-451729, AUDIT-WORKING-963715).

Off-ramp (NOT gold) — routed for adjudication.

- **Acyclicity holds only for the time-unrolled execution graph, not a state-folded policy graph — make the scope explicit.** Gemini (AUDIT-WORKING-193847): the proof (“a cycle would require $\tau_1 < \dots < \tau_i$, impossible for a real time index”) is valid when nodes are *time-indexed events*, but if the strategy DAG were folded over *states of the world* (e.g. “at home” recurring at τ_1 and τ_3), cycles are physical and common — a finite-state-machine / general-MDP transition graph is cyclic. The segment already handles this in its “iteration objection resolved” section (“each attempt is a distinct node at a distinct time”), but the suggestion is to state explicitly at Step 1 that Σ_t is a time-unrolled execution graph, *not* a policy-over-states $\pi(S)$ — because rolling Σ_t up into a state-policy loses the DAG structure. Scope-honesty sharpening of an already-defended point, not a refutation.

Appendix R

Derivation: Sector Condition Verification for Edge-Credence Dynamics

Complete verification that specific strategy update dynamics satisfy the sector condition from A, backing the schema proposed in 9.10. Seven propositions are proved across distinct topologies. **B.1 (single edge)**: sector condition holds globally with $\alpha_\Sigma = 1/(n + 1)$, tight. **B.2 (two-edge AND chain, observable intermediate)**: weakest-link result with $\alpha_\Sigma = \min(1/(n_1 + 1), \theta_1/(n_2 + 1))$ — the evidence-starvation effect on the downstream edge attenuates by the upstream success probability. **B.3 (two-edge AND chain, unobservable intermediate)**: per-edge sector condition *fails* (the marginal point estimate is biased), but plan-level tracking recovers persistence with $\alpha_{\Sigma, \text{plan}} = 1/(n_\Phi + 1)$ at the cost of per-edge diagnostic resolution. **B.4 (two-arm OR, ϵ -greedy)**: *exploration-gated* weakest-link — pure greedy *fails* the sector condition; minimum exploration is required (the sufficient-exploration condition from 9.10).

B.5 (sector transfer to plan-confidence error): the sector condition transfers from per-edge credence space to *strategy-plan-confidence error* via the Jacobian — and the transfer is *credit-assignment-free*, requiring only status propagation. **B.6 (L1 augmented DAG, strict-prerequisite common cause)**: persistence transfers with three-way gating (condition-testing rate, evidence-starvation factor, exploration-gating factor) — the formal extension to causally-insufficient DAGs handled by augmentation. **B.7 (L1' mixture form, soft-facilitator common cause)**: persistence holds under *observable* common cause via five-way gating, and is *refuted* under unobservable common cause by the Cramér-Rao floor (the second formal instance of #disc-identifiability-floor).

The overall result: strategic persistence has a complete verified instance set for the canonical topologies. AND-chains exhibit depth-gated (geometric) attenuation; OR-nodes exhibit exploration-gated allocation; L1-augmented DAGs combine condition-testing, evidence-starvation, and exploration-gating. The chapter's claim that strategic persistence is structurally identical to epistemic persistence at the sector-template level is *verified* by these instantiations.

Motivation.

9.10 proposes that the sector-condition mathematics (Props A.1, A.1S, A.2 in A) extends from epistemic to strategic mismatch. The schema identifies three structural conditions (SA1, SA2', SA3) and conjectures the persistence form $\alpha_\Sigma > \rho_\Sigma/R_\Sigma$. This appendix instantiates the schema for concrete Beta-Bernoulli edge dynamics, verifying the conditions case by case. Each proposition establishes the sector parameter α_Σ for a specific DAG topology; the persistence condition then follows by direct substitution into Proposition A.1 of A.

Common Setup.

Throughout, edge credences follow Beta-Bernoulli dynamics: edge k has true success probability $\theta_k \in (0, 1)$, agent belief $p_k \sim \text{Beta}(\alpha_k, \beta_k)$ with point estimate $\hat{p}_k = \alpha_k/n_k$ where $n_k = \alpha_k + \beta_k$, and conjugate update on observing $y_k \in \{0, 1\}$:

Definition (Beta-Bernoulli edge dynamics).

$$\Delta \hat{p}_k = \frac{y_k - \hat{p}_k}{n_k + 1}, \quad \eta_k = \frac{1}{n_k + 1} \quad (\text{R.1})$$

The single-step update matches the gain-based form from 9.4 with $\eta_{\text{edge}} = 1/(n_k + 1)$.

Parameterization note. Propositions B.1–B.7 below are stated in the *moment-parameter* (probability-space) coordinate $\hat{p}_k \in [0, 1]$, consistent with the Beta-Bernoulli conjugate presentation. The log-odds coordinate $\lambda_k = \log(\hat{p}_k/(1 - \hat{p}_k)) \in \mathbb{R}$ is the unique additive-evidence presentation forced by the evidential-additivity axiom (U); the sector-parameter content of each proposition is Fisher-equivalent in either coordinate, so the derivations are retained in moment-parameter form for algebraic tightness ($\eta_k = 1/(n_k + 1)$ is exact in probability space). The log-odds coordinate becomes load-bearing in 9.6 for the continuous-gradient signal function, where the probability-space presentation would exhibit a mechanical domain break.

Strategic mismatch per edge is $\delta_k = \hat{p}_k - \theta_k$. The sector-condition framework (4.2) requires: (SA1) $F(\mathbf{0}) = \mathbf{0}$, (SA2) $\delta^T \mathbf{F}(\delta) \geq \alpha_\Sigma \|\delta\|^2$ within a region $\mathcal{B}_R$.

Regime-adjustment convention. The propositions below state sector parameters α_Σ for the **canonical Regime-A case**, in which every edge's evidence is fully interventional and attributable (identifiability coefficient $t_{ij} = 1$ per 9.5). For edges with $t_{ij} < 1$ (Regime B or C), each sector parameter carries an implicit t_{ij} factor:

$$\alpha_\Sigma^{\text{eff}} = t_{ij} \cdot \alpha_\Sigma^{\text{stated}} \quad (\text{R.2})$$

This commutes through every diagonal sector product in Props B.1–B.4 and B.6 because each edge's contribution to $\delta_\Sigma^T \mathbf{F}$ is scaled by t_{ij} independently: Regime-C edges ($t \approx 0$) contribute essentially nothing to the weakest-link minimum, and Regime-A edges ($t \approx 1$) recover the stated formulas exactly. The persistence condition $\alpha_\Sigma > \rho_\Sigma/R_\Sigma$ becomes $t_{ij} \cdot \alpha_\Sigma^{\text{stated}} > \rho_{\Sigma,ij}/R_{\Sigma,ij}$ per edge. The propositions below state everything in Regime-A form for clarity; the t -adjustment is uniform across cases and does not alter the structural results (evidence starvation, exploration gating, three-way gating) — only their quantitative thresholds per edge.

For Prop B.3 (unobservable intermediate), the regime adjustment does not change the A1 violation result: the bias is structural to the marginal Bayesian update, not to the identifiability quality. The plan-level fallback still recovers the sector condition with an t -adjusted $\alpha_{\Sigma, \text{plan}}$ governed by the aggregate's identifiability.

Proposition B.1: Single Edge ($A \rightarrow G$).

Setup. One action A , one goal G , one edge with credence $\hat{p}$ and true probability θ . Mismatch: $\delta_\Sigma = \hat{p} - \theta$ (scalar).

Statement. Under Beta-Bernoulli updating, the expected correction function satisfies the sector condition globally (all $|\delta_\Sigma| \leq 1$) with:

Derived (Conditional on Beta-Bernoulli model).

$$\alpha_\Sigma = \frac{1}{n + 1} \quad (\text{R.3})$$

The bound is tight.

Proof.

The expected correction. Since $y \sim \text{Bernoulli}(\theta)$:

Derived (Proof Step).

$$\mathbb{E}[\Delta \hat{p}] = \frac{\theta - \hat{p}}{n+1} = -\frac{\delta_\Sigma}{n+1} \quad (\text{R.4})$$

Identifying $F_\Sigma(\delta_\Sigma) = -\mathbb{E}[\Delta \delta_\Sigma] = \delta_\Sigma/(n+1)$:

Derived (Proof Step).

$$\delta_\Sigma \cdot F_\Sigma(\delta_\Sigma) = \frac{\delta_\Sigma^2}{n+1} = \frac{1}{n+1} \|\delta_\Sigma\|^2 \quad (\text{R.5})$$

So $\alpha_\Sigma = 1/(n+1)$. The correction is exactly linear, so the sector bound is tight (not merely a lower bound).

Verification of (SA1). At $\delta_\Sigma = 0$: $F_\Sigma(0) = 0$. Satisfied.

Stochastic ultimate bound. Decomposing $\delta_{\Sigma,t+1} = \delta_\Sigma \cdot n/(n+1) + (y - \theta)/(n+1)$ and computing $\mathbb{E}[\delta_{\Sigma,t+1}^2]$:

Derived (stochastic ultimate bound, single edge).

$$\mathbb{E}[\delta_{\Sigma,t+1}^2] = \delta_\Sigma^2 \cdot \frac{n^2}{(n+1)^2} + \frac{\theta(1-\theta)}{(n+1)^2} \quad (\text{R.6})$$

The expected Lyapunov derivative $\mathbb{E}[\Delta V] < 0$ whenever:

$$|\delta_\Sigma| > \sqrt{\frac{\theta(1-\theta)}{2n+1}} \quad (\text{R.7})$$

The stochastic noise floor decreases as $O(1/\sqrt{n})$ — the standard Bayesian posterior concentration rate. $\square$

Proposition B.2: Two-Edge AND-Chain, Observable Intermediate ($A \rightarrow B \rightarrow G$, B observed).

Setup. Action A leads to intermediate B , then to goal G . Both y_B and y_G are observed. Edge credences $\hat{p}_1, \hat{p}_2$ with true probabilities θ_1, θ_2 . Plan confidence: $\hat{R}_\Sigma = \hat{p}_1 \hat{p}_2$ (AND propagation from 8.2). Mismatch vector: $\delta_\Sigma = (\delta_1, \delta_2)^T$.

The generative model:

$$y_B \sim \text{Bernoulli}(\theta_1), \quad y_G | y_B = 1 \sim \text{Bernoulli}(\theta_2), \quad y_G | y_B = 0 = 0 \quad (\text{R.8})$$

Statement. Under Beta-Bernoulli updating with observable intermediate, the expected correction function is diagonal and satisfies the sector condition globally with:

Derived (Conditional on Beta-Bernoulli model, observable intermediate).

$$\alpha_\Sigma = \min\left(\frac{1}{n_1+1}, \frac{\theta_1}{n_2+1}\right) \quad (\text{R.9})$$

Proof.

Edge 1 is tested every trial: $\Delta \hat{p}_1 = (y_B - \hat{p}_1)/(n_1 + 1)$, giving $\mathbb{E}[\Delta \hat{p}_1] = -\delta_1/(n_1 + 1)$.

Edge 2 is tested only when $y_B = 1$ (probability θ_1). When $y_B = 0$, edge 2 receives no information:

Derived (Proof Step).

$$\mathbb{E}[\Delta \hat{p}_2] = \theta_1 \cdot \frac{\theta_2 - \hat{p}_2}{n_2 + 1} + (1 - \theta_1) \cdot 0 = -\frac{\theta_1 \delta_2}{n_2 + 1} \quad (\text{R.10})$$

The expected correction function is diagonal:

$$\mathbf{F}(\delta_\Sigma) = \begin{pmatrix} \delta_1/(n_1 + 1) \\ \theta_1 \delta_2/(n_2 + 1) \end{pmatrix} \quad (\text{R.11})$$

Verification of (SA1). At $\delta_\Sigma = \mathbf{0}$: $\mathbf{F}(\mathbf{0}) = \mathbf{0}$. Satisfied.

The sector product:

Derived (Proof Step).

$$\delta_\Sigma^T \mathbf{F}(\delta_\Sigma) = \frac{\delta_1^2}{n_1 + 1} + \frac{\theta_1 \delta_2^2}{n_2 + 1} \geq \min\left(\frac{1}{n_1 + 1}, \frac{\theta_1}{n_2 + 1}\right) (\delta_1^2 + \delta_2^2) \quad (\text{R.12})$$

The minimum is achieved by whichever edge has the lower effective correction rate. $\square$

The evidence-starvation effect. Edge 2's effective correction rate $\theta_1/(n_2 + 1)$ is attenuated by the upstream success probability θ_1 . When θ_1 is small, edge 2 is starved of evidence — it is tested rarely, learns slowly, and becomes the weakest link. This gives a precise mechanism for 8.1: downstream edges in AND-chains are harder to calibrate, with correction rate proportional to the product of all upstream success probabilities.

Generalization to depth d . For a chain of depth d with all intermediates observable, edge k 's effective correction rate is $\alpha_k = \prod_{j=1}^{k-1} \theta_j/(n_k + 1)$, giving:

Derived (depth-d chain sector parameter).

$$\alpha_\Sigma = \min_{k=1}^d \frac{\prod_{j=1}^{k-1} \theta_j}{n_k + 1} \quad (\text{R.13})$$

Proposition B.3: Two-Edge AND-Chain, Unobservable Intermediate ($A \rightarrow B \rightarrow G$, B not observed).

Setup. Same DAG as Proposition B.2, but only $y_G \in \{0, 1\}$ is observed. The agent cannot distinguish “ $A \rightarrow B$ failed” from “ $B \rightarrow G$ failed” when $y_G = 0$.

Statement. (a) The per-edge sector condition fails: the marginal Bayesian update violates (SA1) with systematic bias $O(1/n)$. (b) Plan-level tracking (treating $\hat{\Phi} = \hat{p}_1 \hat{p}_2$ as a single Beta over the composite probability $\Phi = \theta_1 \theta_2$) recovers the sector condition with:

Derived (Conditional on Beta-Bernoulli model, unobservable intermediate).

$$\alpha_{\Sigma, \text{plan}} = \frac{1}{n_{\Phi} + 1} \quad (\text{R.14})$$

Proof of (a): Per-edge (SA1) violation.

The joint posterior after success ($y_G = 1$) factors cleanly: $\alpha_k \rightarrow \alpha_k + 1$ for both edges. The joint posterior after failure ($y_G = 0$) includes the factor $(1 - \theta_1 \theta_2)$, which does not decompose into separate functions of θ_1 and θ_2 . Failure introduces posterior correlation.

The exact marginal Bayesian update for edge 1 on failure is:

Derived (Proof Step).

$$\Delta \hat{p}_1^{(\text{fail})} = -\frac{\hat{p}_1 q_1}{n_1 + 1} \quad (\text{R.15})$$

where $q_1 = (1 - \hat{p}_1)/(1 - \hat{p}_1 \hat{p}_2)$ is the posterior probability that edge 1 caused the failure (the proportional-blame weight). The expected update, combining success (probability $\theta_1 \theta_2$) and failure (probability $1 - \theta_1 \theta_2$):

Derived (Proof Step).

$$\mathbb{E}[\Delta \hat{p}_1] = \theta_1 \theta_2 \cdot \frac{1 - \hat{p}_1}{n_1 + 1} - (1 - \theta_1 \theta_2) \cdot \frac{\hat{p}_1 q_1}{n_1 + 1} \quad (\text{R.16})$$

Evaluating at truth ($\hat{p}_k = \theta_k$):

Derived (SA1 violation magnitude).

$$\mathbb{E}[\Delta \hat{p}_1]_{\hat{p}=\theta} = -\frac{\theta_1(1 - \theta_1)(1 - \theta_2)}{n_1 + 1} \neq 0 \quad (\text{R.17})$$

The bias is negative (pushes credence below truth), of magnitude $O(1/n)$, and arises because success always increments α_1 by a full unit while failure increments β_1 by the fractional weight $q_1 < 1$. The success case over-credits edge 1 because the agent cannot separate edge 1’s contribution from edge 2’s. By symmetry, $\hat{p}_2$ is also biased downward. $\square$

Proof of (b): Plan-level recovery.

If the agent tracks $\hat{\Phi}$ directly as a Beta posterior over $\Phi = \theta_1 \theta_2$ using $y_G \sim \text{Bernoulli}(\Phi)$, the setup reduces to a single-edge problem (Proposition B.1) with mismatch $\delta_{\Phi} = \hat{\Phi} - \Phi$ and sector parameter $\alpha_{\Sigma, \text{plan}} = 1/(n_{\Phi} + 1)$. All single-edge results apply. $\square$

Tradeoff. Per-edge tracking provides diagnostic resolution (which edge is failing?) but yields biased, non-sector-conforming updates when the intermediate is unobservable. Plan-level tracking sacrifices diagnostic resolution but

provides clean sector-condition satisfaction. This is a concrete instance of 9.3: unobservable nodes freeze per-edge learning, making the aggregate the correct unit of analysis.

Non-identifiability. From y_G observations alone, only $\Phi = \theta_1 \theta_2$ is identifiable. The per-edge decomposition (θ_1, θ_2) is underdetermined — any pair with the same product produces identical observations. Per-edge identification requires per-edge observability (or structural constraints that break the symmetry).

Proposition B.4: Two-Arm OR-Node ($A_1, A_2 \rightarrow G$, ε -greedy).

Setup. Goal G reachable via two alternative actions A_1, A_2 (an OR-node from 8.2). True probabilities θ_1, θ_2 . Agent selects one action per trial: with probability $1 - \varepsilon$ the greedy choice (higher $\hat{p}_k$), with probability ε the other. Mismatch vector: $\delta_\Sigma = (\delta_1, \delta_2)^T$.

Statement. (a) Under a greedy policy ($\varepsilon = 0$), the sector condition fails for the full mismatch vector. (b) Under ε -greedy action selection with $\varepsilon > 0$, the sector condition is satisfied with:

Derived (Conditional on Beta-Bernoulli model, ε -greedy OR-node).

$$\alpha_\Sigma = \min\left(\frac{1 - \varepsilon}{n_1 + 1}, \frac{\varepsilon}{n_2 + 1}\right) \quad (\text{R.18})$$

where arm 1 is the greedy choice.

Proof of (a): Greedy failure.

Under greedy selection, only the chosen arm updates. Assuming $\hat{p}_1 > \hat{p}_2$ (arm 1 is greedy):

$$\mathbf{F}(\delta_\Sigma) = \begin{pmatrix} \delta_1/(n_1 + 1) \\ 0 \end{pmatrix} \quad (\text{R.19})$$

The sector product $\delta_\Sigma^T \mathbf{F} = \delta_1^2/(n_1 + 1)$. For the sector condition to hold, we need $\delta_1^2/(n_1 + 1) \geq \alpha_\Sigma(\delta_1^2 + \delta_2^2)$. Setting $\delta_1 = 0$: the LHS is zero but $\alpha_\Sigma \delta_2^2 > 0$ for $\delta_2 \neq 0$. No $\alpha_\Sigma > 0$ works. $\square$

Proof of (b): ε -greedy recovery.

The expected correction (arm 1 greedy):

Derived (Proof Step).

$$\mathbb{E}[\mathbf{F}(\delta_\Sigma)] = \begin{pmatrix} (1 - \varepsilon) \delta_1/(n_1 + 1) \\ \varepsilon \delta_2/(n_2 + 1) \end{pmatrix} \quad (\text{R.20})$$

The sector product:

Derived (Proof Step).

$$\delta_\Sigma^T \mathbf{F} = \frac{(1 - \varepsilon) \delta_1^2}{n_1 + 1} + \frac{\varepsilon \delta_2^2}{n_2 + 1} \geq \min\left(\frac{1 - \varepsilon}{n_1 + 1}, \frac{\varepsilon}{n_2 + 1}\right)(\delta_1^2 + \delta_2^2) \quad (\text{R.21})$$

Verification of (SA1). At $\delta_\Sigma = \mathbf{0}$: $\mathbf{F}(\mathbf{0}) = \mathbf{0}$. Satisfied. $\square$

Optimal exploration rate. Maximizing α_Σ over ε by equalizing the two terms:

Derived (optimal exploration rate).

$$\varepsilon^* = \frac{n_2 + 1}{n_1 + n_2 + 2}, \quad \alpha_\Sigma^* = \frac{1}{n_1 + n_2 + 2} \quad (\text{R.22})$$

For equal experience ($n_1 = n_2$): $\varepsilon^* = 1/2$ (equal allocation). For unequal experience, ε^* allocates more trials to the arm with higher n (lower gain) to equalize correction rates.

The k -arm generalization. For k OR-alternatives with ε -uniform exploration (probability $\varepsilon/(k-1)$ per non-greedy arm):

Derived (k -arm OR-node sector parameter).

$$\alpha_\Sigma = \min\left(\frac{1 - \varepsilon}{n_{\text{greedy}} + 1}, \min_{j \neq \text{greedy}} \frac{\varepsilon/(k-1)}{n_j + 1}\right) \quad (\text{R.23})$$

With optimal equal-rate exploration and equal experience: $\alpha_\Sigma = 1/(k(n+1))$. The sector parameter scales as $1/k$ — each additional alternative dilutes correction capacity.

Condition (SA3). For the sector condition to be satisfiable at all, the exploration rate must satisfy:

Derived (minimum exploration rate for persistence).

$$\varepsilon > \frac{\rho_\Sigma(n_{\max} + 1)}{R_\Sigma} \quad (\text{R.24})$$

A purely greedy agent ($\varepsilon = 0$) violates the sector condition because the unchosen alternative's mismatch grows without correction. This is the structural justification for (SA3) in 9.10: OR-nodes require deliberate exploration to maintain strategic persistence.

Persistence Conditions.

Substituting each proposition's α_Σ into the persistence threshold $\alpha_\Sigma > \rho_\Sigma/R_\Sigma$ from Proposition A.1 of A:

Single edge (B.1):

Derived (single-edge persistence threshold).

$$\frac{1}{n+1} > \frac{\rho_\Sigma}{R_\Sigma} \quad \Longleftrightarrow \quad n < \frac{R_\Sigma}{\rho_\Sigma} - 1 \quad (\text{R.25})$$

The critical experience level $n^* = R_\Sigma/\rho_\Sigma - 1$ is the point at which gain collapse occurs: the posterior has concentrated so much that the agent can no longer track environmental drift.

Two-edge observable (B.2):

Derived (two-edge persistence threshold, observable).

$$\min\left(\frac{1}{n_1 + 1}, \frac{\theta_1}{n_2 + 1}\right) > \frac{\rho_\Sigma}{R_\Sigma} \quad (\text{R.26})$$

This decomposes into two independent conditions: $n_1 < R_\Sigma/\rho_\Sigma - 1$ and $n_2 < \theta_1 R_\Sigma/\rho_\Sigma - 1$. The downstream edge has a stricter bound, tightened by θ_1 . Gain collapse hits downstream edges first.

Two-edge unobservable (B.3):

$$\frac{1}{n_\Phi + 1} > \frac{\rho_\Sigma}{R_\Sigma} \iff n_\Phi < \frac{R_\Sigma}{\rho_\Sigma} - 1 \quad (\text{R.27})$$

Same form as single-edge, applied to the plan-level aggregate.

Two-arm OR, ε -greedy (B.4):

Derived (OR-node persistence threshold).

$$\min\left(\frac{1 - \varepsilon}{n_1 + 1}, \frac{\varepsilon}{n_2 + 1}\right) > \frac{\rho_\Sigma}{R_\Sigma} \quad (\text{R.28})$$

With optimal exploration and equal experience ($n_1 = n_2 = n$):

$$n < \frac{R_\Sigma}{2\rho_\Sigma} - 1 \quad (\text{R.29})$$

The critical experience is halved relative to the single-edge case. Each additional OR-alternative further reduces the critical level by $1/k$.

Proposition B.6: Mixed AND/OR with Common-Cause Node (L1 Augmented DAG).

Setup. Goal G reachable via two alternative actions A_1, A_2 (an OR-node), but both actions require a shared condition C (a common-cause node). The L1 DAG factors the common cause above the OR-structure: $G = \text{AND}(C, G_{\text{sub}})$ where $G_{\text{sub}} = \text{OR}(A_1, A_2)$. Three leaves: condition C with true probability θ_C , actions A_1, A_2 with true conditional probabilities $\theta_{1|C}, \theta_{2|C}$. Agent selects ε -greedy among A_1, A_2 (only testable when $C = 1$). Mismatch vector: $\delta_\Sigma = (\delta_C, \delta_{A_1}, \delta_{A_2})^T$.

Plan confidence: $\hat{R}_\Sigma = \hat{p}_C \cdot [1 - (1 - \hat{p}_{A_1})(1 - \hat{p}_{A_2})]$. Reference value: $\Phi = \theta_C \cdot [1 - (1 - \theta_{1|C})(1 - \theta_{2|C})]$. At truth, $\hat{R}_\Sigma = \Phi = \text{actual plan success probability (L1 is unbiased)}$.

Statement. Under Beta-Bernoulli updating with observable condition C and ε -greedy path selection, the expected correction function satisfies the sector condition globally with:

Derived (Conditional on Beta-Bernoulli model, L1 augmented DAG, observable common cause C , ε -greedy action selection).

$$\alpha_\Sigma = \min\left(\frac{1}{n_C + 1}, \frac{\theta_C(1 - \varepsilon)}{n_{A_1} + 1}, \frac{\theta_C\varepsilon}{n_{A_2} + 1}\right) \quad (\text{R.30})$$

where A_1 is the greedy choice.

Proof.

Edge C (condition leaf, observed every trial):

$$\mathbb{E}[\Delta \hat{p}_C] = \frac{\theta_C - \hat{p}_C}{n_C + 1} = -\frac{\delta_C}{n_C + 1} \quad (\text{R.31})$$

Edge A_1 (action leaf, tested when $C = 1$ and path 1 selected — probability $\theta_C(1 - \varepsilon)$):

$$\mathbb{E}[\Delta \hat{p}_{A_1}] = -\frac{\theta_C(1 - \varepsilon) \delta_{A_1}}{n_{A_1} + 1} \quad (\text{R.32})$$

Edge A_2 (action leaf, tested when $C = 1$ and path 2 selected — probability $\theta_C \varepsilon$):

$$\mathbb{E}[\Delta \hat{p}_{A_2}] = -\frac{\theta_C \varepsilon \delta_{A_2}}{n_{A_2} + 1} \quad (\text{R.33})$$

The expected correction function is diagonal:

$$\mathbf{F}(\delta_\Sigma) = \begin{pmatrix} \delta_C/(n_C + 1) \\ \theta_C(1 - \varepsilon) \delta_{A_1}/(n_{A_1} + 1) \\ \theta_C \varepsilon \delta_{A_2}/(n_{A_2} + 1) \end{pmatrix} \quad (\text{R.34})$$

Verification of (SA1). $\mathbf{F}(\mathbf{0}) = \mathbf{0}$. ✓

Sector product:

$$\delta_\Sigma^T \mathbf{F} = \frac{\delta_C^2}{n_C + 1} + \frac{\theta_C(1 - \varepsilon) \delta_{A_1}^2}{n_{A_1} + 1} + \frac{\theta_C \varepsilon \delta_{A_2}^2}{n_{A_2} + 1} \geq \alpha_\Sigma \|\delta_\Sigma\|^2 \quad \square \quad (\text{R.35})$$

Three-way gating. This is the first mixed AND/OR case. The sector parameter α_Σ is determined by three independent gating mechanisms: 1. **Condition testing** ($1/(n_C + 1)$): the common cause is the easiest component to calibrate — tested every trial, same as B.1. 2. **Evidence starvation** (θ_C factor on action terms): both action edges are gated by the condition probability, same mechanism as B.2. 3. **Exploration gating** ($(1 - \varepsilon)$ and ε factors): the two OR-alternatives compete for test opportunities, same mechanism as B.4.

The bottleneck is typically the explore action behind the condition: $\theta_C \varepsilon / (n_{A_2} + 1)$.

B.5b applies. The Jacobian $\mathbf{J} = (s_{G_{\text{sub}}}, \hat{p}_C(1 - \hat{p}_{A_2}), \hat{p}_C(1 - \hat{p}_{A_1}))^T$ is non-negative (monotone AND/OR). Corrections are componentwise (each leaf updates independently). Therefore $\alpha_s = \alpha_c = \alpha_\Sigma$ — the strategy-plan-confidence error inherits the sector condition with no penalty.

Optimal exploration. Equalizing the action terms: $\varepsilon^* = (n_{A_2} + 1)/(n_{A_1} + n_{A_2} + 2)$, giving $\alpha_\Sigma^* = \min(1/(n_C + 1), \theta_C/(n_{A_1} + n_{A_2} + 2))$.

L0 comparison. An L0 model of the same scenario (no C node, marginal probabilities $\theta_k = \theta_C \cdot \theta_{k|C}$) has $\alpha_\Sigma^{L0} = \min((1 - \varepsilon)/(n_1 + 1), \varepsilon/(n_2 + 1))$ and reference value $\Phi^{L0} = 1 - (1 - \theta_1)(1 - \theta_2)$. The L0 sector parameter is *higher* (no θ_C attenuation) but Φ^{L0} systematically overestimates actual plan success. L1 persistence is harder (lower α_Σ) but

calibration is honest (Φ^{L1} matches reality). The tradeoff: L0 is easier to maintain but calibrates to a biased target; L1 is harder to maintain but calibrates to truth.

The L1 construction principle. The common-cause node C must be factored *above* the OR-structure whose children it correlates. A naive construction (adding C as a parent of both A_1 and A_2 while keeping them as OR-siblings) does not fix the overestimation because the OR-propagation formula treats siblings as marginally independent. Correct L1 construction makes C an AND-prerequisite above the OR-structure, so the OR-children are conditionally independent given C . Full worked example with quantitative L0/L1 comparison: #example-L1.

Summary of Results.

Table R.1

Case	Topology	α_Σ	SAI	SA3	Weakest link
B.1 Single edge	$A \rightarrow G$	$1/(n+1)$	Yes	N/A	—
B.2 Two-edge, B observable	$A \rightarrow B \rightarrow G$ (AND)	$\min(1/(n_1+1), \theta_1/(n_2+1))$	Yes	N/A	Depth-gated (evidence starvation)
B.3 Two-edge, B unobservable	$A \rightarrow B \rightarrow G$ (AND)	$1/(n_\Phi+1)$ (plan-level)	Per-edge: No ($O(1/n)$ bias); Plan: Yes	N/A	Credit-assignment collapse
B.4 Two-arm OR, ε -greedy	$A_1, A_2 \rightarrow G$ (OR)	$\min((1-\varepsilon)/(n_1+1), \varepsilon/(n_2+1))$	Yes	Required	Exploration-gated
B.5a Credence $\rightarrow$ value (linear)	Any	$\alpha_s = \alpha_c$ (Jacobian cancels)	—	—	None (exact transfer)
B.5b Credence $\rightarrow$ value (nonlinear, component-wise)	Any	$\alpha_s = \alpha_c$ ($J_k \geq 0$ preserves bound)	—	—	None (exact transfer)
B.5b Credence $\rightarrow$ value (coupled)	Any	$\alpha_s \geq \alpha_c/\kappa(J)^2$	—	—	Inter-edge coupling
B.6 L1 augmented, mixed AND/OR	$G = \text{AND}(C, \text{OR}(A_1, A_2))$	$\min(1/(n_C+1), \theta_C(1-\varepsilon)/(n_{A_1}+1), \theta_C\varepsilon/(n_{A_2}+1))$	Yes	Required	Three-way gating
B.7 L1' mixture form, observable C	$\hat{P}_\Sigma^{L1'} = \hat{\theta}_C P_\Sigma(\hat{\mathbf{p}}_{ C}) + (1-\hat{\theta}_C) P_\Sigma(\hat{\mathbf{p}}_{ \neg C})$	$\min(1/(n_C+1), \min_j \theta_C \pi_{j C}/(n_{j C}+1), \min_j (1-\theta_C) \pi_{j \neg C}/(n_{j \neg C}+1))$	Yes	Required and facilitator monotonicity	Five-way gating; refuted under unobservable C (Cramér-Rao floor)

Structural results across cases:

- **Gain as sector parameter.** In every case, α_Σ is determined by the edge update gain $\eta_k = 1/(n_k + 1)$, modulated by observability (θ_1 attenuation in B.2), credit-assignment (plan-level collapse in B.3), action-selection policy (ε in B.4), or a combination of these (B.6). The structural parallel between epistemic and strategic persistence is not an analogy — it is a mathematical identity at the sector-framework level.
- **Evidence starvation (AND-chains).** Downstream edges in AND-chains have correction rates attenuated by the product of upstream success probabilities. Depth increases fragility: α_Σ decays exponentially with chain depth.
- **Exploration gating (OR-nodes).** Unchosen OR-alternatives receive zero correction. The sector condition requires deliberate exploration (SA3) at a rate exceeding $\rho_\Sigma(n_{\max} + 1)/R_\Sigma$. Pure greedy policies fail.
- **Three-way gating (L1 augmented).** When a common cause gates OR-alternatives, calibration faces condition testing, evidence starvation, and exploration gating simultaneously (B.6). The bottleneck is the least-explored action behind the rarest condition.
- **Five-way gating (L1' mixture, observable C).** When the common cause is a *soft facilitator* ($\theta_{j|-C} > 0$) and is observable per trial, calibration faces five gating mechanisms simultaneously (B.7): condition testing, evidence starvation on each conditional branch ($C = 1$ and $C = 0$), and exploration gating on each branch. The bottleneck is whichever branch is rarer, modulated by exploration on that branch. **Identifiability obstruction under unobservable C:** without direct C -observation, the mixture parameters are non-identifiable from a single observation channel — the per-trial Fisher information matrix is rank 1 rather than rank $2K + 1$, so by the Cramér-Rao bound (Cramér 1946, *Mathematical Methods of Statistics*, Princeton University Press) no unbiased online estimator on the joint conditional vector admits $\alpha > 0$. The agent must augment C -observability, run multi-child joint observations, or fall back to plan-level (L0-on-marginals) tracking.
- **Gain-collapse threshold.** In all cases, increasing experience eventually violates persistence when the environment drifts. The critical experience level n^* is inversely proportional to the drift rate ρ_Σ .
- **L1 tradeoff.** L1 augmented DAGs have lower α_Σ (harder to persist) but honest calibration (Φ^{L1} matches reality). L0 is easier to maintain but calibrates to a biased target. The choice depends on whether accuracy or trackability is the binding constraint.

Proposition B.7: L1' Mixture-Form Sector Transfer (Observable Common Cause).

Where Prop B.6 handles a strict-prerequisite common cause via an AND-prerequisite construction (L1), this proposition handles a *soft-facilitator* common cause via the mixture form L1' (8.3, Correlation Hierarchy). It is the L1' analog of B.6, generalizing three-way gating to five-way gating across two parallel conditional branches.

Setup. L1' (mixture-form) DAG with binary common cause C , true prior $\theta_C \in (0, 1)$, gating a sub-plan with conditional structures $G_{j|C}$ (when $C = 1$) and $G_{j|-C}$ (when $C = 0$). Each affected child j carries two conditional credences $\hat{p}_{j|C}$ and $\hat{p}_{j|-C}$. Plan confidence:

$$\hat{p}_\Sigma^{L1'} = \hat{\theta}_C \cdot R_\Sigma(\hat{p}_{j|C}) + (1 - \hat{\theta}_C) \cdot R_\Sigma(\hat{p}_{j|-C}) \quad (R.36)$$

The agent observes C -state directly each trial (i.e., C is treated as an observable condition leaf, in parallel with B.6). Per trial, the agent (i) updates $\hat{\theta}_C$ from observed C via Beta-Bernoulli; (ii) on $C = 1$ trials, attempts a child j via the active selection policy $\pi_{j|C}$ and updates $\hat{p}_{j|C}$ from y_j ; (iii) on $C = 0$ trials, attempts a child j via $\pi_{j|-C}$ and updates $\hat{p}_{j|-C}$ from y_j .

Mismatch state: $\xi = (\xi_C, \{\xi_{j|C}\}, \{\xi_{j|-C}\})$ with $\xi_C = \theta_C - \hat{\theta}_C$, $\xi_{j|s} = \theta_{j|s} - \hat{p}_{j|s}$.

Statement.

Under Beta-Bernoulli updating with observable common cause, componentwise updates per conditional branch, and *facilitator monotonicity* ($R_\Sigma(G|_C) \geq R_\Sigma(G|_{\neg C})$), the expected correction function on the joint state satisfies the sector condition globally with:

Derived (Conditional on Beta-Bernoulli model, observable common cause C, componentwise conditional-branch update, facilitator monotonicity).

$$\alpha_{L1'} = \min\left(\frac{1}{n_C + 1}, \min_{j \in \mathcal{J}_C} \frac{\theta_C \pi_{j|C}}{n_{j|C} + 1}, \min_{j \in \mathcal{J}_{\neg C}} \frac{(1 - \theta_C) \pi_{j|\neg C}}{n_{j|\neg C} + 1}\right) \quad (\text{R.37})$$

where $\mathcal{J}_s$ is the set of children tested on the $C = s$ branch and $\pi_{j|s}$ is the action-selection probability for child j on that branch.

Proof. *Edge C (condition leaf, observable every trial).* Standard Beta-Bernoulli on direct C -observation:

$$\mathbb{E}[\Delta \hat{\theta}_C] = \frac{\theta_C - \hat{\theta}_C}{n_C + 1} = -\frac{\xi_C}{n_C + 1} \quad (\text{R.38})$$

so $F_C(\xi) = \xi_C / (n_C + 1)$.

Edge $p_{j|C}$ (conditional credence, $C = 1$ branch). Tested when $C = 1$ (probability θ_C) and the agent's $C = 1$ -conditional policy selects j (probability $\pi_{j|C}$):

$$\mathbb{E}[\Delta \hat{p}_{j|C}] = \theta_C \pi_{j|C} \cdot \frac{\theta_{j|C} - \hat{p}_{j|C}}{n_{j|C} + 1} = -\frac{\theta_C \pi_{j|C} \xi_{j|C}}{n_{j|C} + 1} \quad (\text{R.39})$$

so $F_{j|C}(\xi) = \theta_C \pi_{j|C} \xi_{j|C} / (n_{j|C} + 1)$.

Edge $p_{j|\neg C}$ (conditional credence, $C = 0$ branch). Symmetric: tested with probability $(1 - \theta_C) \pi_{j|\neg C}$:

$$F_{j|\neg C}(\xi) = \frac{(1 - \theta_C) \pi_{j|\neg C} \xi_{j|\neg C}}{n_{j|\neg C} + 1} \quad (\text{R.40})$$

Sector product. The correction function is diagonal in ξ with positive entries:

$$\xi^T \mathbf{F}(\xi) = \frac{\xi_C^2}{n_C + 1} + \sum_{j \in \mathcal{J}_C} \frac{\theta_C \pi_{j|C} \xi_{j|C}^2}{n_{j|C} + 1} + \sum_{j \in \mathcal{J}_{\neg C}} \frac{(1 - \theta_C) \pi_{j|\neg C} \xi_{j|\neg C}^2}{n_{j|\neg C} + 1} \geq \alpha_{L1'} \cdot \|\xi\|^2 \quad (\text{R.41})$$

with $\alpha_{L1'}$ as stated.

SA1 check. $\mathbf{F}(\mathbf{0}) = \mathbf{0}$ at truth ($\xi = 0$). $\checkmark$

Globality. The Beta-Bernoulli updates are linear in ξ on $[-1, 1]^{2K+1}$, so the diagonal Jacobian computed at truth is the exact $\mathbf{J}_F$ throughout. The sector inequality holds globally with the same $\alpha_{L1'}$.

B.5b bridge to plan value. Jacobian of $\hat{P}_\Sigma^{L1'}$ with respect to the joint state:

$$\mathbf{J}_{P_\Sigma} = \begin{pmatrix} R_\Sigma(\hat{\mathbf{p}}_{|C}) - R_\Sigma(\hat{\mathbf{p}}_{|\neg C}) \\ \hat{\theta}_C \cdot \nabla_{\hat{\mathbf{p}}_{|C}} R_\Sigma(\hat{\mathbf{p}}_{|C}) \\ (1 - \hat{\theta}_C) \cdot \nabla_{\hat{\mathbf{p}}_{|\neg C}} R_\Sigma(\hat{\mathbf{p}}_{|\neg C}) \end{pmatrix} \quad (\text{R.42})$$

By facilitator monotonicity, the first entry is ≥ 0 . The other entries are ≥ 0 by AND/OR monotonicity on each conditional sub-DAG. So $\mathbf{J}_{P_\Sigma} \geq 0$ componentwise. By the componentwise nonlinear case of B.5b, the sector condition transfers losslessly: $\alpha_s = \alpha_c = \alpha_{L1}$. $\square$

Five-Way Gating. This generalizes B.6's three-way gating to five gating mechanisms, one per term in the min:

1. **Condition testing** ($1/(n_C + 1)$): the common cause is the easiest component to calibrate — directly observed every trial.
2. **Evidence starvation, $C = 1$ branch** (θ_C factor): conditional edges on the $C = 1$ branch are gated by the prior frequency of $C = 1$.
3. **Exploration gating, $C = 1$ branch** ($\pi_{j|C}$ factor): OR-alternatives within the $C = 1$ branch compete for test opportunities.
4. **Evidence starvation, $C = 0$ branch** ($1 - \theta_C$ factor): conditional edges on the $C = 0$ branch are gated by the prior frequency of $C = 0$.
5. **Exploration gating, $C = 0$ branch** ($\pi_{j|\neg C}$ factor): OR-alternatives within the $C = 0$ branch.

The bottleneck is typically whichever of θ_C and $1 - \theta_C$ is smaller (the rare branch sees fewer trials), modulated by exploration on that branch.

Reduction to B.6 (Strict-Prerequisite Limit). Setting $\theta_{j|\neg C} \rightarrow 0$ collapses the $\neg C$ branch (the agent learns the $\neg C$ conditionals are zero with no remaining variance). The per-edge state for the $\neg C$ branch drops out of the sector inequality, and the formula reduces to B.6's three-way gating with $\alpha_\Sigma = \min(1/(n_C + 1), \theta_C(1 - \varepsilon)/(n_{A_1} + 1), \theta_C\varepsilon/(n_{A_2} + 1))$. $\checkmark$

Refuted Under Unobservable C (Mixture Identifiability Obstruction).

When C is not observable, the agent can only update the joint state via online soft EM on the marginal y_j . Computing the Fisher information of the mixture model $\mu_j = \theta_C \theta_{j|C} + (1 - \theta_C) \theta_{j|\neg C}$ at truth, the score vector for parameters $\phi = (\theta_C, p_{j|C}, p_{j|\neg C})$ admits the rank-1 factorization:

Derived (cramer-rao-floor, from Fisher information of mixture model + Cramér-Rao bound).

$$\mathcal{F}(\phi) = \frac{1}{\mu_j(1 - \mu_j)} uu^T, \quad u = (\Delta_j, \theta_C, 1 - \theta_C) \quad (\text{R.43})$$

where $\Delta_j = p_{j|C} - p_{j|\neg C}$ is the *separability gap*. The matrix is rank 1 for any non-degenerate mixture; the two-dimensional null space corresponds to perturbations along the indeterminacy manifold $\{\hat{\phi} : \hat{\theta}_C \hat{p}_{|C} + (1 - \hat{\theta}_C) \hat{p}_{|\neg C} = \mu_j\}$ — directions unobservable from a single binary signal.

Since the soft-EM step at truth equals (up to the $1/(n + 1)$ scaling) the natural-gradient ascent on the expected log-likelihood, the Jacobian of the correction function equals $\mathcal{F}/(n + 1)$, also rank 1. Therefore $\lambda_{\min}(\mathbf{J}_F) = 0$, and **no SAI-preserving update on the joint conditional vector admits a sector parameter $\alpha > 0$ under SUB-B**. This is a Cramér-Rao bound, not a defect of any particular update rule: any unbiased online estimator must respect the bound, which is infinite in the unidentifiable directions.

Repair routes. When C is unobservable, the agent must:

(i) **Augment C -observability** (recover B.7 above) — instrument secondary signals that identify C per trial, transforming the problem from refuted to globally derived. (ii) **Run $K \geq 2$ children jointly under the same C -realization** — when the joint Fisher matrix reaches rank $2K + 1$, the mixture is identifiable; local sector condition holds with $\alpha \geq \sigma_{\min}(\mathcal{F}_{\text{joint}}) / \max_k(n_k + 1)$. Strong structural requirement (joint observation typically unavailable in standard sequential strategy execution). (iii) **Fall back to plan-level tracking on the marginal $\hat{\mu}_j$** — recovers B.1's $\alpha = 1/(n_\mu + 1)$ on the scalar marginal but loses the per-conditional decomposition (equivalent to L0-on-marginals).

Connection to the Identifiability-Floor Pattern. The B.7-vs-refuted-unobservable- C structure is one of two no-go theorems in AAT's strategy layer. The other is the on-policy detection no-go in 9.2, which prohibits L0-vs-L1 distinction from purely on-policy data via the causal hierarchy theorem. Both are structural impossibility results derived from external information-theoretic theorems (Bareinboim et al. 2022 Causal Hierarchy Theorem; Cramér-Rao bound on the Fisher information). Both strengthen related machinery — the on-policy no-go strengthens 7.3; the mixture no-go strengthens the case for *observability-as-information-augmentation*. See 5.1 for the meta-pattern.

Proposition B.5: Bridge from Credence Error to Value Residuals.

The propositions above verify the sector condition for the credence-error mismatch state $\delta_c = (\hat{p}_k - \theta_k)$. This proposition extends the result to **strategy-plan-confidence error** — the scalar difference between the agent's strategy-plan-confidence score and the independence-model plan value evaluated at true edge parameters. This is a natural plan-level mismatch that is distinct from (but related to) the per-edge value-increment residuals $\delta_{\text{strategic}}$ defined in 9.1.

Relationship to $\delta_{\text{strategic}}$. 9.1 defines $\delta_{\text{strategic}}$ as an L^2 aggregation of per-edge value-increment residuals — a quantity that requires credit assignment to compute. Plan-confidence error δ_s defined here is a *different* quantity: it is the error in the DAG's aggregate self-assessment, computable from status propagation without credit assignment. The two are related (both measure strategy-reality mismatch) but not identical. This proposition shows the sector condition holds for strategy-plan-confidence error; extending it to $\delta_{\text{strategic}}$ directly would require the credit-assignment machinery discussed in 9.6.

Setup. The plan value $\hat{P}_\Sigma = R_\Sigma(\mathbf{p})$ is a function of the edge credence vector $\mathbf{p} = (p_1, \dots, p_m)$. The **independence-model reference value** is $\Phi = R_\Sigma(\theta)$ — the AND/OR propagation formula evaluated at the true edge parameters θ . Note: Φ is NOT the actual probability of plan success when edge outcomes are correlated (8.3, edge-independence caveat). It is the best the independence model can do with perfect edge knowledge. Under correlated failure, the actual plan success probability may be lower than Φ ; the gap is a model-class limitation of the AND/OR DAG representation, not an estimation error. What δ_s tracks is calibration *within the independence model*, not calibration to reality. Define:

- **Credence-error mismatch:** $\delta_c = \mathbf{p} - \theta \in \mathbb{R}^m$
- **Plan-confidence error:** $\delta_s = \hat{P}_\Sigma - \Phi = R_\Sigma(\mathbf{p}) - R_\Sigma(\theta)$ (scalar)

The gradient of plan value with respect to credences is $\mathbf{J} = \nabla_{\mathbf{p}} R_\Sigma \in \mathbb{R}^m$ (a vector, since R_Σ is scalar). By first-order Taylor expansion:

$$\delta_s \approx \mathbf{J}^T \delta_c \tag{R.44}$$

B.5a: Linear Correction (Beta-Bernoulli).

When the correction in credence space is linear — $\mathbf{F}_c(\delta_c) = \eta \odot \delta_c$ where $\eta_k = 1/(n_k + 1)$ and $\odot$ is elementwise product — consider the sector ratio directly. Since $\delta_s = \mathbf{J}^T \delta_c$ and $F_s = \mathbf{J}^T \mathbf{F}_c$:

Derived (value-residual sector condition, linear case).

$$\delta_s \cdot F_s = (\mathbf{J}^T \delta_c)^T (\mathbf{J}^T \mathbf{F}_c) = \delta_c^T \mathbf{J} \mathbf{J}^T \mathbf{F}_c \quad (\text{R.45})$$

For the linear case $\mathbf{F}_c = \eta_{\min} \delta_c$ (using the worst-case uniform gain for the bound):

$$\delta_s \cdot F_s = \eta_{\min} \cdot \delta_c^T \mathbf{J} \mathbf{J}^T \delta_c = \eta_{\min} \cdot \|\mathbf{J}^T \delta_c\|^2 = \eta_{\min} \cdot \delta_s^2 \quad (\text{R.46})$$

Therefore:

$$\frac{\delta_s \cdot F_s}{\delta_s^2} = \eta_{\min} \quad (\text{R.47})$$

The Jacobian cancels. The sector parameter in value-residual space equals the sector parameter in credence space, regardless of DAG structure:

$$\alpha_s = \alpha_c = \eta_{\min} = \min_k \frac{1}{n_k + 1} \quad (\text{R.48})$$

This holds because $\mathbf{J} \mathbf{J}^T$ appears in both numerator and denominator of the sector ratio, and the linear correction commutes with the Jacobian mapping.

Interpretation. The Jacobian $\mathbf{J}$ determines *which* credence errors matter most for plan value (high-sensitivity edges contribute more to value residuals). But it does not change the *rate* at which those errors are corrected, because the correction is proportional to the error regardless of the error's plan-value impact.

B.5b: Nonlinear Componentwise Correction.

When the correction is nonlinear but **componentwise** — each edge corrects independently, with

Derived (from per-component sector condition + Jacobian non-negativity).

$$(\mathbf{F}_c)_k \text{ depending only on } (\delta_c)_k \quad (\text{R.49})$$

— and the plan-value Jacobian is **non-negative** ($J_k \geq 0$ for all k), the transfer is lossless regardless of nonlinearity.

The per-component sector condition (from Props B.1-B.4): for each edge k ,

$$(\delta_c)_k \cdot (\mathbf{F}_c)_k \geq \alpha_c \cdot (\delta_c)_k^2 \quad (\text{R.50})$$

which gives

$$(\mathbf{F}_c)_k / (\delta_c)_k \geq \alpha_c \quad (\text{R.51})$$

(the correction-to-mismatch ratio exceeds α_c for each edge independently).

Jacobian non-negativity holds for all well-formed AND/OR DAGs: increasing any edge credence p_k never decreases plan value R_Σ (monotonicity of AND/OR propagation). Therefore $J_k = \partial R_\Sigma / \partial p_k \geq 0$ for all k .

The transfer. Since $J_k \geq 0$ preserves the inequality direction:

$$J_k \cdot (\mathbf{F}_c)_k \geq J_k \cdot \alpha_c \cdot (\delta_c)_k \quad \text{for each } k \quad (\text{R.52})$$

Summing:

$$\mathbf{J}^T \mathbf{F}_c = \sum_k J_k (\mathbf{F}_c)_k \geq \alpha_c \sum_k J_k (\delta_c)_k = \alpha_c \cdot \mathbf{J}^T \delta_c \quad (\text{R.53})$$

Therefore $F_s \geq \alpha_c \cdot \delta_s$, giving:

$$\alpha_s = \alpha_c \quad (\text{R.54})$$

No condition-number penalty. The Jacobian's anisotropy is irrelevant when the correction is componentwise and the Jacobian is non-negative. The non-negativity ensures that every edge's correction contributes in the right direction to the value correction, regardless of how unevenly the Jacobian weights different edges.

When does componentwise correction hold? For all cases in Props B.1, B.2, and B.4 — the standard Beta-Bernoulli update on each edge, including the stochastic (nonlinear) single-step realization. The correction IS nonlinear (the actual per-step update $(y - \hat{p})/(n + 1)$ is random, not proportional to δ_k), but it is componentwise (each edge updates from its own observation independently).

When does it fail? When edge corrections are coupled — correcting one edge changes another's correction direction. This occurs for: - **Unobservable intermediates (B.3) under marginal Bayesian updates:** The marginal Bayesian update couples edge posteriors. The correction for edge 2 depends on the joint posterior over both edges, not just on $(\delta_c)_2$ alone. *And — critically — Prop B.3(a) showed that this scheme violates SA1 with $O(1/n)$ bias: there is no valid α_c for the marginal Bayesian update on coupled edges in the first place.* The per-edge bound must therefore be re-established before the Jacobian transfer can be applied. - **Shared observations:** If testing one arm gives partial information about another (correlated arms), corrections couple across edges — with an SA1 status that depends on the specific attribution scheme used to project the shared observation onto per-edge updates.

The coupled-case Cauchy-Schwarz bound presupposes a valid α_c . For the bound $\alpha_s \geq \alpha_c / \kappa(\mathbf{J})^2$ to apply meaningfully, the per-edge sector condition must hold — i.e., the attribution scheme must satisfy SA1. Marginal Bayesian updates on coupled edges do *not* satisfy this precondition (Prop B.3(a)). The resolution is to replace marginal Bayesian with **gradient-based attribution** (Prop B.5d below), which recovers SA1 and yields a valid per-edge sector parameter that admits the Jacobian transfer with the κ^2 penalty.

Scope (coupled-case-applicability).

Per-edge vs. plan-level routes through coupling. Prop B.3(b) already established an alternative route for the two-edge unobservable case: tracking at the plan level ($\hat{\Phi} = \hat{p}_1 \hat{p}_2$) with a single Beta posterior recovers SA1 directly, yielding $\alpha_\Sigma = 1/(n_\Phi + 1)$. The plan-level route works *without* needing the Jacobian bridge because it sidesteps per-edge attribution entirely. The gradient-based per-edge route works *through* the Jacobian bridge, preserving per-edge credences at the cost of the κ^2 penalty. Both routes give valid sector parameters; the choice depends on whether per-edge diagnostics or aggregate efficiency matters more.

The refined picture: nonlinearity is not the problem; inter-edge coupling combined with the attribution scheme is. A nonlinear but componentwise correction transfers losslessly through any non-negative Jacobian. A coupled correction under a scheme that preserves SA1 (gradient-based) incurs the condition-number penalty because the coupling

can redirect correction away from the value-relevant direction. A coupled correction under a scheme that violates SA1 (marginal Bayesian) has no valid α_c to transfer at all and must be replaced before B.5 applies.

Proposition B.5d: Gradient-Based Attribution Restores SA1 on Coupled Edges. **Setup.** Same two-edge AND topology as B.3: edges with true probabilities θ_1, θ_2 , credences $\hat{p}_1, \hat{p}_2$, plan probability $\hat{R}_\Sigma = \hat{p}_1 \hat{p}_2$, observable plan outcome $y_G \in \{0, 1\}$ with $\mathbb{E}[y_G] = \Phi = \theta_1 \theta_2$. Intermediate is unobservable.

Gradient-based per-edge update. Distribute plan-level surprise proportionally to the Jacobian component:

Formulation (gradient-attribution).

$$\Delta \hat{p}_k = \frac{1}{n_k + 1} \cdot J_k \cdot (y_G - \hat{R}_\Sigma), \quad J_k = \frac{\partial \hat{R}_\Sigma}{\partial \hat{p}_k} \quad (\text{R.55})$$

For the two-edge AND case, $J_1 = \hat{p}_2, J_2 = \hat{p}_1$.

Verification of SA1. At truth ($\hat{p}_k = \theta_k$):

$$\mathbb{E}[\Delta \hat{p}_k]_{\hat{\mathbf{p}}=\theta} = \frac{J_k}{n_k + 1} \cdot \mathbb{E}[y_G - \hat{R}_\Sigma]_{\hat{\mathbf{p}}=\theta} \quad (\text{R.56})$$

The bracket evaluates as: $\mathbb{E}[y_G] = \Phi$ by definition of y_G as a Bernoulli outcome with mean Φ ; and $\hat{R}_\Sigma|_{\hat{\mathbf{p}}=\theta} = \theta_1 \theta_2 = \Phi$. Therefore $\mathbb{E}[y_G - \hat{R}_\Sigma]_{\hat{\mathbf{p}}=\theta} = 0$, and:

$$\mathbb{E}[\Delta \hat{p}_k]_{\hat{\mathbf{p}}=\theta} = 0 \quad \text{for all } k. \quad (\text{SA1 satisfied}) \quad (\text{R.57})$$

The contrast with marginal Bayesian is precise: the marginal scheme suffers from a Φ -vs- θ_k mismatch (the update treats each edge's marginal likelihood rather than the joint), producing the $O(1/n)$ bias derived in Prop B.3(a). The gradient scheme distributes surprise via the Jacobian, which aligns the per-edge update direction with the plan-level gradient — and this gradient is zero at truth by construction.

Sector parameter (SA2'). With the gradient-based scheme, the per-edge correction is:

$$F_k(\delta) = \frac{J_k}{n_k + 1} \cdot (\Phi - \hat{R}_\Sigma) \approx \frac{J_k}{n_k + 1} \cdot \mathbf{J}^T \delta \quad (\text{R.58})$$

where the approximation uses the first-order expansion $\hat{R}_\Sigma - \Phi \approx \mathbf{J}^T \delta$ valid near truth. The per-edge sector bound is:

$$\delta^T \mathbf{F}(\delta) \approx \sum_k \frac{J_k^2}{n_k + 1} \cdot (\mathbf{J}^T \delta)^2 / \|\delta\| \geq \frac{\sigma_{\min}(\mathbf{J})^2}{\max_k(n_k + 1)} \|\delta\|^2 \quad (\text{R.59})$$

giving:

$$\alpha_c^{\text{grad}} \geq \frac{\sigma_{\min}(\mathbf{J})^2}{\max_k(n_k + 1)} \quad (\text{R.60})$$

Transfer through B.5c. Applying the coupled-case Cauchy-Schwarz bound:

$$\alpha_s \geq \frac{\alpha_c^{\text{grad}}}{\kappa(\mathbf{J})^2} = \frac{\sigma_{\min}(\mathbf{J})^2 / \max_k(n_k + 1)}{\sigma_{\max}(\mathbf{J})^2 / \sigma_{\min}(\mathbf{J})^2} = \frac{\sigma_{\min}(\mathbf{J})^4}{\sigma_{\max}(\mathbf{J})^2 \cdot \max_k(n_k + 1)} \quad (\text{R.61})$$

The $\sigma_{\min}(\mathbf{J})^4 / \sigma_{\max}(\mathbf{J})^2$ factor is the full cost of coupled per-edge attribution — a steep penalty compared to the componentwise case. The plan-level route (B.3(b)) avoids this penalty entirely by tracking Φ as a single quantity with $\alpha_\Sigma = 1/(n_\Phi + 1)$, at the cost of losing per-edge diagnostics.

Why not proportional-blame or other schemes? Proportional-blame distributes surprise proportionally to current credences: $\Delta \hat{p}_k \propto \hat{p}_k \cdot (y_G - \hat{R}_\Sigma) / \sum_j \hat{p}_j$. This violates SA1 because the proportionality factor $\hat{p}_k / \sum_j \hat{p}_j$ is not the gradient direction, so the distribution over edges does not align with the plan-level surprise’s actual attribution. Only schemes whose per-edge update aligns with the Jacobian direction preserve SA1 for coupled edges; gradient-based is the canonical such scheme.

Claim. For coupled edges with unobservable intermediates, gradient-based attribution is the minimal scheme that (a) satisfies SA1, (b) yields a well-defined per-edge sector parameter, (c) admits the Cauchy-Schwarz Jacobian transfer, and (d) permits per-edge diagnostics. Its $\sigma_{\min}^4 / \sigma_{\max}^2$ scaling is inherent to the coupled case, not an artifact of the scheme. $\square$

B.5c: Implications. Four regimes for the credence-to-value bridge, indexed by correction type and attribution scheme:

Table R.2

Correction type	Edge coupling	Attribution	SA1	Transfer	α_s
Linear	Any	—	Yes	Exact (Jacobian cancels)	α_c
Nonlinear, componentwise	Independent edges	Per-edge Bayesian	Yes	Exact ($J_k \geq 0$ preserves bound)	α_c
Nonlinear, coupled	Coupled edges	Gradient-based (Jacobian distribution)	Yes	$\kappa(\mathbf{J})^2$ penalty	$\alpha_c^{\text{grad}} / \kappa^2$
Nonlinear, coupled	Coupled edges	Marginal Bayesian	No (O(1/n) bias)	— (Jacobian bridge does not apply)	plan-level alternative (8.3)

Componentwise correction closes the gap for all verified cases. Props B.1, B.2, and B.4 all use componentwise edge updates with non-negative Jacobian. The operational diagnostic $\delta_{\text{strategic}}$ inherits the sector condition from credence error with no penalty — even for the nonlinear (stochastic) single-step realization. The derivation validates the operational mismatch, not merely a surrogate.

Inter-edge coupling — not nonlinearity — is the fragility. The κ^2 penalty arises only when correcting one edge changes another’s correction direction (B.3’s unobservable case, correlated observations). Deep or unbalanced DAGs have high $\kappa(\mathbf{J})$ when corrections couple, but the coupling, not the depth or imbalance per se, is what degrades the persistence guarantee.

Connection to credit assignment. The Jacobian $\mathbf{J} = \nabla_{\mathbf{p}} R_\Sigma$ is computable from the DAG’s status propagation formulas — it does not require solving the credit-assignment problem. The credit-assignment problem is about *decomposing observed value changes into per-edge contributions* (needed for the update rule). The Jacobian bridge is about *transferring a per-edge sector condition to the value-residual space* (needed for the persistence guarantee). These are different problems: the bridge works even when credit assignment is unsolved, because it only requires the sensitivity structure, not the causal attribution.

What Is Derived vs. What Is Chosen.

The segment carries derivations at three distinct strengths: (i) exact sector-parameter algebra for concrete Beta-Bernoulli topologies (B.1, B.2, B.4, B.6, B.7 under observable C + facilitator monotonicity), (ii) two structural refutations that are *themselves* derivations of impossibility (B.3(a) SA1 bias; B.7 Cramér-Rao floor under unobservable C), and (iii) formulation choices (Beta-Bernoulli dynamics, componentwise conditional-branch updates, the particular ε -greedy policy family). The dividing line runs between the Beta-Bernoulli *model* (chosen) and the sector-parameter *algebra conditional on that model* (proved). B.7's positive transfer requires two load-bearing conditions — observability of C and facilitator monotonicity $R_\Sigma(G \mid C) \geq R_\Sigma(G \mid \neg C)$ — either of which failing disqualifies the transfer.

Table R.3

<i>Property</i>	<i>Source</i>	<i>Strength</i>
Beta-Bernoulli edge dynamics ($\Delta \hat{p}_k = (y_k - \hat{p}_k)/(n_k + 1)$)	Conjugate-prior model; matches 9.4	Formulation choice
Regime-A adjustment: $\alpha_\Sigma^{\text{eff}} = t_{ij} \cdot \alpha_\Sigma^{\text{stated}}$	Diagonal commutation across Props B.1–B.4, B.6; 9.5	Derived
B.1 Single-edge sector parameter $\alpha_\Sigma = 1/(n + 1)$	Expected Beta-Bernoulli update + sector product algebra	Proved (tight, conditional on Beta-Bernoulli)
B.1 Stochastic ultimate bound $ \delta > \sqrt{\theta(1 - \theta)/(2n + 1)}$	Second-moment recursion + expected-Lyapunov decrement	Derived
B.2 Observable AND-chain sector parameter $\min(1/(n_1 + 1), \theta_1/(n_2 + 1))$	Diagonal correction function; sector product algebra	Proved (conditional on Beta-Bernoulli + observable intermediate)
B.2 Depth- d generalization $\alpha_\Sigma = \min_k \prod_{j < k} \theta_j/(n_k + 1)$	Iteration of B.2 pattern under observable intermediates	Derived
B.3(a) Per-edge SA1 violation under marginal Bayesian attribution ($O(1/n)$ bias)	Exact marginal update + evaluation at truth	Derived as refutation (no valid α_c exists)
B.3(b) Plan-level recovery $\alpha_{\Sigma, \text{plan}} = 1/(n_\Phi + 1)$	Reduction of aggregate $\hat{\Phi}$ tracking to B.1	Derived (reduction)
B.3 Non-identifiability of (θ_1, θ_2) from y_G alone	Product $\Phi = \theta_1 \theta_2$ symmetry in observation channel	Proved
B.4(a) Pure-greedy OR ($\varepsilon = 0$) violates sector condition	Counterexample at $\delta_1 = 0, \delta_2 \neq 0$	Derived as refutation
B.4(b) ε -greedy sector parameter $\min((1 - \varepsilon)/(n_1 + 1), \varepsilon/(n_2 + 1))$	Policy-weighted correction; sector product algebra	Proved (conditional on Beta-Bernoulli + ε -greedy)
B.4 Optimal exploration $\varepsilon^* = (n_2 + 1)/(n_1 + n_2 + 2)$	Term-equalization in min	Derived
B.4 Minimum exploration rate (SA3 requirement) $\varepsilon > \rho_\Sigma(n_{\max} + 1)/R_\Sigma$	Substitution into persistence threshold $\alpha_\Sigma > \rho_\Sigma/R_\Sigma$	Derived
B.5a Linear credence-to-value transfer $\alpha_s = \alpha_c$ (Jacobian cancels)	Algebraic cancellation of $\mathbf{J}\mathbf{J}^T$	Proved (exact, DAG-structure independent)
B.5b Componentwise nonlinear transfer $\alpha_s = \alpha_c$	Per-component sector bound + monotone AND/OR ($\mathbf{J} \geq 0$)	Proved (conditional on componentwise update + non-negative Jacobian)
B.5b Coupled-edge Cauchy-Schwarz bound $\alpha_s \geq \alpha_c/\kappa(\mathbf{J})^2$	Cauchy-Schwarz on $\mathbf{J}^T \mathbf{F}_c$ and $\mathbf{J}^T \hat{\delta}_c$	Derived (conditional on Jacobian regularity and an SA1-preserving attribution scheme)
B.5c Four-regime classification of credence-to-value transfer	Combination of linear/nonlinear $\times$ componentwise/coupled $\times$ attribution type	Derived
B.5d Gradient-based attribution restores SA1 on coupled edges	Jacobian-weighted surprise distribution; $\mathbb{E}[y_G - \hat{p}_\Sigma]_{ \theta} = 0$	Derived

Property	Source	Strength
B.5d Sector parameter under gradient attribution $\alpha_c^{\text{grad}} \geq \sigma_{\min}(\mathbf{J})^2 / \max_k(n_k + 1)$	First-order expansion near truth + singular-value analysis	Derived
B.5d Gradient-based as <i>minimal</i> SA1-preserving scheme for coupled edges	Argument excluding proportional-blame and other schemes	Derived qualitatively (not a uniqueness theorem)
B.6 L1 augmented-DAG sector parameter (three-way gating: condition / starvation / exploration)	Diagonal correction function over $(\delta_C, \delta_{A_1}, \delta_{A_2})$; sector product	Proved (conditional on Beta-Bernoulli + observable C + ε -greedy)
B.6 L0-vs-L1 comparison: L1 is lower α_Σ but honestly calibrated	Direct Φ^{L0} vs Φ^{L1} computation at truth	Derived
B.6 L1-construction principle (common cause factored <i>above</i> OR-structure)	Requirement for conditional independence of OR-siblings given C	Derived
B.7 L1' mixture-form sector parameter (five-way gating)	Diagonal correction function over $(\xi_C, \{\xi_{j C}\}, \{\xi_{j \neg C}\})$; sector product	Derived (conditional on Beta-Bernoulli + observable C + componentwise conditional-branch update + <i>facilitator monotonicity</i> $P_\Sigma(G C) \geq P_\Sigma(G \neg C)$)
B.7 Componentwise update per conditional branch	Matches Beta-Bernoulli on the branch actually executed each trial	Formulation choice
B.7 Facilitator-monotonicity condition $P_\Sigma(G C) \geq P_\Sigma(G \neg C)$	Required so the first Jacobian entry of $\hat{P}_\Sigma^{L1'}$ is non-negative, enabling the lossless B.5b transfer	Load-bearing scope condition
B.7 Reduction to B.6 under strict-prerequisite limit ($\theta_{j \neg C} \rightarrow 0$)	Collapse of $\neg C$ -branch terms in the sector inequality	Verification (algebraic consistency)
B.7 Refutation under unobservable C (single-channel Fisher rank 1)	Cramér-Rao bound applied to rank-1 Fisher $\mathcal{F}(\phi) = uu^T / [\mu_j(1 - \mu_j)]$	Proved by external theorem (Cramér-Rao bound) — no unbiased online estimator admits $\alpha > 0$
B.7 Repair route (i): augment C -observability	Recovers the positive B.7 derivation	Formulation / scope choice (maps to 7.3)
B.7 Repair route (ii): $K \geq 2$ joint children under same C -realization	Fisher matrix rank reaches $2K + 1$; local sector condition	Derived (sketch; applicability narrow)
B.7 Repair route (iii): plan-level fallback on marginal $\hat{\mu}_j$	Reduction to B.1 on scalar marginal	Derived (at cost of per-conditional decomposition)
Persistence thresholds (substitution of each α_Σ into $\alpha_\Sigma > \rho_\Sigma / R_\Sigma$)	Proposition A.1 of A	Derived (each)
Gain-collapse threshold $n^* \propto R_\Sigma / \rho_\Sigma$ common across cases	Direct inversion of persistence condition	Derived
Expected-value (rather than almost-sure) sector analysis	Pedagogical simplification; a.s. convergence via standard Bayesian consistency	Formulation choice (scope)
Constant- α assumption with experience-discounting (λ) stabilization	Instantaneous persistence check + forgetting-factor analysis	Derived (instantaneous only)

Epistemic Status.

Conditional on the Beta-Bernoulli model. Propositions B.1, B.2, B.4, and B.6 are *derived*: the sector-condition verification is exact algebra under the stated generative model. The persistence conditions follow by direct application of Proposition A.1 (A), which is independently established. Proposition B.3(a) (the SA1 violation) is also derived; B.3(b) reduces to B.1 by construction. Proposition B.5a (the credence-to-value bridge for linear correction) is *exact* — the Jacobian cancellation is algebraic, independent of DAG structure. Proposition B.5b (the componentwise nonlinear transfer) is *exact under non-negative Jacobian* — algebraic, no condition-number penalty. The coupled-edge case of

B.5b is *conditional on (i) Jacobian regularity* — the condition-number bound requires $\sigma_{\min}(\mathbf{J}) > 0$, which holds for non-degenerate DAGs — *and (ii) an SA1-preserving attribution scheme* (gradient-based, Prop B.5d); marginal Bayesian attribution on coupled edges does not admit the Cauchy-Schwarz bound because SA1 fails upstream. Proposition B.5d (gradient-based attribution) is *derived* — SA1 verification is algebraic and the sector-parameter bound follows from standard singular-value analysis of the Jacobian.

All results use the *expected-value* sector condition. A full stochastic treatment (Foster-Lyapunov or supermartingale convergence) would give probability bounds rather than expected-value bounds. The expected-value analysis gives the correct asymptotic behavior (posterior concentration at $O(1/\sqrt{n})$) but does not prove almost-sure convergence. For the stationary Beta-Bernoulli case, almost-sure convergence is guaranteed by the standard Bayesian consistency theorem independently of this derivation.

The time-varying α_{Σ} issue remains: since n_k increases with each observation, α_{Σ} decreases over time. The sector-condition framework assumes constant α . The results here give an *instantaneous* persistence check at the current experience level, not a trajectory guarantee. Experience discounting (exponential forgetting with factor λ) stabilizes α_{Σ} at the exact value $(1 - \lambda)/(2 - \lambda)$ from $1/(n_{\text{eff}} + 1)$ with $n_{\text{eff}} = 1/(1 - \lambda)$ — with the simpler $\approx 1 - \lambda$ as the slow-forgetting asymptote ($\lambda \rightarrow 1$) — yielding the persistence requirement $(1 - \lambda)/(2 - \lambda) > \rho_{\Sigma}/R_{\Sigma}$ (full treatment in 9.10).

What remains open:

1. *General DAG topology.* **Partially resolved.** Proposition B.6 verifies the first mixed AND/OR case (L1 augmented DAG with common-cause node). The three-way gating structure (condition testing $\times$ evidence starvation $\times$ exploration gating) appears to be the general pattern. Remaining: arbitrary mixed AND/OR DAGs with multiple common causes and deeper nesting.
2. *Continuous outcomes.* The Beta-Bernoulli model gives conjugate, closed-form updates. Non-conjugate cases (continuous signals, partial observability) require approximate inference, and the sector condition must be verified for the approximation.
3. *Modified sector condition for biased correction.* **Resolved via Prop B.5d (gradient-based attribution).** The $O(1/n)$ bias in the unobservable case (B.3a) under marginal Bayesian updates is eliminated by switching to gradient-based attribution, which satisfies SA1 exactly at truth. Per-edge results are recovered under the gradient scheme at the $\sigma_{\min}^4/\sigma_{\max}^2$ scaling. Alternative route (plan-level tracking, B.3(b)) remains available without the condition-number penalty.
4. *Correlated edges (L1/L2 scope).* **Resolved.** Proposition B.6 verifies the sector condition for L1 (strict-prerequisite augmented DAG); Proposition B.7 verifies it for L1' (soft-facilitator mixture form, observable common cause) with five-way gating. The unobservable- C single-channel case is *refuted* by the Cramér-Rao floor (mixture identifiability obstruction; see B.7 §"Refuted Under Unobservable C ") — not merely "open" but structurally impossible without additional information augmentation. Practical paths under unobservable C : augment C -observability, run multi-child joint observations (sketch only, narrow applicability), or fall back to plan-level tracking on the marginal. Strategies requiring full L2 conditioning over k unobservable common causes still cost $O(2^k)$.
5. *Adaptive exploration.* Proposition B.4 uses fixed ε . Adaptive strategies (UCB, Thompson sampling) allocate exploration based on current uncertainty and should yield tighter sector bounds.

Discussion.

The sector parameter is the edge update gain. Across all four cases, α_Σ is a function of $\eta_k = 1/(n_k + 1)$ — the same quantity that governs epistemic persistence (3.8, 3.6). Strategic persistence and epistemic persistence are governed by the same mathematical machinery, not merely structurally parallel. This validates the schema proposed in 9.10 at the level of concrete dynamics.

AND vs. OR: structural vs. behavioral fragility. The evidence-starvation effect (B.2) is structural — it depends on the DAG topology and upstream reliability, not on the agent’s policy. The exploration-gating effect (B.4) is behavioral — it depends on the action-selection policy. AND-node persistence is improved by increasing upstream reliability (making early steps more likely to succeed). OR-node persistence is improved by increasing exploration (testing alternatives more frequently). The distinction maps to the familiar depth-vs-breadth tradeoff in planning.

Gain collapse as the dominant failure mode. In every drifting-environment case, the persistence threshold takes the form $n < c/\rho_\Sigma$ for some constant c . Accumulated experience suppresses the update gain, eventually making the agent unable to track environmental change. This gives the gain-collapse phenomenon (3.6) a precise quantitative threshold for each topology.

The credence-to-value bridge decouples persistence from credit assignment. Proposition B.5 shows that strategic persistence (in value-residual space) can be established from per-edge sector conditions (in credence space) without solving the credit-assignment problem. The bridge requires only the value-sensitivity Jacobian $\mathbf{J} = \nabla_{\mathbf{p}} P_\Sigma$, which is computable from the DAG’s status propagation formulas in $O(|V| + |E|)$ time. Credit assignment — attributing observed value changes to specific edges — is needed for the *update rule* (computing the signal function in 9.4), not for the *persistence guarantee*. This means the persistence analysis can proceed even while the update rule remains incompletely specified.

Inter-edge coupling is the true fragility, not nonlinearity or DAG depth. Proposition B.5b shows that nonlinear componentwise corrections transfer losslessly through any non-negative Jacobian — the $\kappa(\mathbf{J})^2$ penalty arises only when corrections are *coupled* across edges. This refines the earlier qualitative argument from 8.1: the structural pressure toward shallow, balanced strategies is not about depth per se but about the coupling that deep or complex structures tend to introduce. A deep AND-chain with fully observable intermediates (B.2) has componentwise corrections and transfers losslessly despite high depth. The same chain with unobservable intermediates (B.3) has coupled corrections and incurs the penalty. The distinction is observability-mediated coupling, not topological complexity.

Appendix S

Derivation: Hard Ceiling on Strategic-Persistence Reachability under Exponential Forgetting

The strategic-persistence schema's instantaneous persistence form $\alpha_\Sigma > \rho_\Sigma/R_\Sigma$ (9.10), instantiated under Beta-Bernoulli edge dynamics with exponential forgetting at rate $\lambda \in (0, 1)$, has a sharp structural cap at $\rho_\Sigma/R_\Sigma = 1/2$: no λ satisfies the prerequisite when strategic disturbance reaches half the strategic reserve. The schema's reachable persistence region is exactly the open half-plane $\rho_\Sigma < R_\Sigma/2$.

Formal Expression.

Setup. Per R Prop B.1, the Beta-Bernoulli edge update with n accumulated pseudo-counts gives sector parameter $\alpha_\Sigma = 1/(n + 1)$. Without forgetting, $n \rightarrow \infty$ and $\alpha_\Sigma \rightarrow 0$ for every edge asymptotically — so the schema's instantaneous persistence form eventually fails under any positive disturbance rate. *Exponential forgetting* with discount $\lambda \in (0, 1)$ replaces the raw update with the discounted recurrence:

Definition (Discounted Beta-Bernoulli Update).

$$\alpha_k \mapsto \lambda \alpha_k + y_k, \quad \beta_k \mapsto \lambda \beta_k + (1 - y_k) \quad (\text{S.1})$$

following the standard adaptive-control / online-learning treatment (Ljung 1987).

Proposition C.1 (Steady-State Sector Parameter).

Under the discounted Beta-Bernoulli recurrence, the steady-state sector parameter at the fixed point $\hat{p} = \theta$ is:

Derived (Conditional on Beta-Bernoulli edges + exponential forgetting).

$$\alpha_\Sigma^{\text{ss}} = \frac{1 - \lambda}{2 - \lambda} \quad (\text{S.2})$$

Proof. At a fixed point in expectation, the discounted recurrence stabilizes when each pseudo-count equals its own contribution rate divided by the dissipation:

$$\alpha^* = \lambda \alpha^* + \theta \implies \alpha^* = \frac{\theta}{1-\lambda}, \quad \beta^* = \lambda \beta^* + (1-\theta) \implies \beta^* = \frac{1-\theta}{1-\lambda} \quad (\text{S.3})$$

The effective sample size is the sum:

$$n_{\text{eff}} = \alpha^* + \beta^* = \frac{1}{1-\lambda} \quad (\text{S.4})$$

Substituting into Prop B.1's $\alpha_\Sigma = 1/(n+1)$ at $n = n_{\text{eff}}$:

$$\alpha_\Sigma^{\text{ss}} = \frac{1}{n_{\text{eff}} + 1} = \frac{1}{1/(1-\lambda) + 1} = \frac{1-\lambda}{1+(1-\lambda)} = \frac{1-\lambda}{2-\lambda} \quad \square \quad (\text{S.5})$$

The slow-forgetting linear form $\alpha_\Sigma^{\text{ss}} \approx 1-\lambda$ is the asymptotic expansion as $\lambda \rightarrow 1$ (the regime where the $1/(2-\lambda)$ damping factor approaches unity). Outside the slow-forgetting limit the linear form overstates the steady-state $\alpha_\Sigma^{\text{ss}}$ — at $\lambda = 0.5$ by 50% (1/3 exact vs 1/2 linear); at $\lambda = 0.9$ by 10% (1/11 exact vs 1/10 linear). The exact form is operationally the threshold to check against; the linear form is unsafe outside its asymptotic regime.

Proposition C.2 (Hard-Ceiling No-Go).

The forgetting prerequisite for the schema's trajectory guarantee under exponential forgetting is:

Derived (Conditional on Prop C.1 + the schema's persistence form).

$$\frac{1-\lambda}{2-\lambda} > \frac{\rho_\Sigma}{R_\Sigma} \quad (\text{S.6})$$

This inequality is unsatisfiable for any $\lambda \in [0, 1]$ when $\rho_\Sigma \geq R_\Sigma/2$. The supremum is sharp:

$$\sup_{\lambda \in [0,1]} \alpha_\Sigma^{\text{ss}}(\lambda) = \frac{1}{2} \quad (\text{S.7})$$

achieved as $\lambda \rightarrow 0^+$ (maximally aggressive forgetting — no memory). The schema's reachable persistence region in (ρ_Σ, R_Σ) -space is exactly the open half-plane $\rho_\Sigma < R_\Sigma/2$.

Proof: Let $x = \rho_\Sigma/R_\Sigma$ with $x \in [0, 1]$. Solving for the threshold λ at which the inequality becomes equality:

$$(1-\lambda) = x(2-\lambda) \implies \lambda(x-1) = 2x-1 \implies \lambda = \frac{2x-1}{x-1} = \frac{1-2x}{1-x} \quad (\text{S.8})$$

For $\lambda \in (0, 1)$ to admit a strict solution to the inequality, we need $\lambda > 0$ at the threshold. With $x < 1$ (denominator positive), this requires $1-2x > 0$, i.e., $x < 1/2$.

At $x = 1/2$ exactly: the threshold $\lambda = 0$, giving $\alpha_\Sigma^{\text{ss}}(0) = (1-0)/(2-0) = 1/2 = x$ — equality, not strict satisfaction. For $x > 1/2$: no $\lambda \in [0, 1]$ satisfies the strict inequality, since even maximally aggressive forgetting ($\lambda \rightarrow 0^+$) only achieves $\alpha_\Sigma^{\text{ss}} \rightarrow 1/2$. The supremum is exactly 1/2 at $\lambda = 0$. $\square$

Equivalently: the schema's reachable set under any exponential-forgetting design is exactly $\{(\rho_\Sigma, R_\Sigma) : \rho_\Sigma < R_\Sigma/2\}$. The boundary $\rho_\Sigma = R_\Sigma/2$ is not in the reachable region (strict satisfaction required), and the half-plane above the boundary is entirely outside.

Epistemic Status.

Exact under the named conditions: (i) Beta-Bernoulli edge dynamics (R Prop B.1), (ii) exponential forgetting with $\lambda \in (0, 1)$, (iii) the schema's persistence form per 9.10 and the underlying template-instantiation conditions (T1)–(T3) of G. Both Propositions C.1 and C.2 are algebraically exact within these conditions — no slack, no approximation. The “exact” tier reflects that uncertainty resides in whether the conditions hold (which mismatch state, which topology, whether the forgetting mechanism is in fact exponential vs another discounting scheme), not in the derivation.

The hard ceiling is *structural*: it is independent of λ choice and independent of which topology among the schema's verified instances (single-edge, two-edge AND observable/unobservable, two-arm OR with ε -greedy, mixed AND/OR with common cause). It is a *class-level cap on reachability* — not a refutation of the schema's persistence form, but a refinement of which (ρ_Σ, R_Σ) regimes the schema can be satisfied in.

Discussion.

Why this no-go matters. The hard ceiling *refines*, not *refutes*, the schema. The schema's persistence form $\alpha_\Sigma > \rho_\Sigma/R_\Sigma$ remains the right form; what the ceiling establishes is that under exponential forgetting, the form is *unreachable* across half the (ρ_Σ, R_Σ) -space. Above the ceiling, structural persistence of Σ_t is not available to this architecture — no tuning of λ can recover it. This is the canonical analog at the strategic layer of 4.8's epistemic-layer trigger: when parametric adaptation cannot keep up, structural change is required.

Boundary behavior. At $\rho_\Sigma/R_\Sigma = 1/2$ exactly, the supremum is achieved at $\lambda = 0$ — the *no-memory* regime where every observation fully replaces the prior. This corresponds to effective sample size $n_{\text{eff}} = 1$ (only the most recent observation counts), giving $\alpha_\Sigma = 1/2$ in Prop B.1's $1/(n + 1)$. The reachable region's open-versus-closed character at the boundary is itself informative: persistence requires *strict* satisfaction, so the boundary $\rho_\Sigma = R_\Sigma/2$ is not in the reachable region even with $\lambda = 0$.

Connection to the broader gain-decay class. The hard ceiling sharpens to a class-level no-go on a wider class of update mechanisms — see NeurIPS 2026 Paper 2 ($\mathcal{A}_{\text{decay}}$ structural-class theorem) for the lift from Beta-Bernoulli with forgetting to all gain-decay updates: every mechanism whose effective gain decays to zero with accumulated experience universally violates the persistence threshold for any positive disturbance rate, and finite-gain mechanisms face *bidirectional* thresholds (this lower bound from the schema's persistence form, plus an upper bound from noise blow-up). The hard ceiling at $1/2$ is the Beta-Bernoulli + exponential-forgetting instance of this broader structural pattern.

Strict-form independent-verify note. Per 5 and the project's verification discipline, the canonical strict-form check on this derivation is a fresh-mathematician re-derivation of Propositions C.1 and C.2 (algebra above) plus independent reading of R Prop B.1 (the source the C.1 derivation reduces to). This is owed before any further status elevation; the present “exact” tier rests on the algebra's elementariness plus the conditional structure being clean.

Findings.

The Hard Ceiling at $\rho_\Sigma = R_\Sigma/2$ (Class-Level Reachability Cap). **Brief:** Under the schema's Beta-Bernoulli + exponential-forgetting instantiation, $\sup_{\lambda \in [0,1]} \alpha_\Sigma^{\text{ss}} = 1/2$, achieved as $\lambda \rightarrow 0^+$. The schema's reachable persistence region in (ρ_Σ, R_Σ) -space is exactly the open half-plane $\rho_\Sigma < R_\Sigma/2$. When the strategic disturbance rate reaches half the strategic reserve, no choice of forgetting rate satisfies the persistence prerequisite — structural persistence of Σ_t is unavailable to this architecture above the ceiling.

Impact: Identifies a sharp, λ -independent cap on what exponential forgetting can achieve. The cap is not an artifact of tuning; it is structural — a class-level no-go on the gain-decay-plus-forgetting class of strategic update mechanisms. For an agent whose strategic-disturbance-to-reserve ratio approaches $1/2$, no engineering of the forgetting mechanism can restore the schema's trajectory guarantee. The hard ceiling closes the design space: above it, the agent must change *something else* — the update mechanism class (away from gain-decay-plus-forgetting toward constant-step finite-gain,

per the NeurIPS Paper 2 $\mathcal{A}_{\text{decay}}$ structural-class theorem), the mismatch state, or the topology. This is the canonical strategic-layer analog of 4.8’s epistemic-layer trigger.

Novelty Claim: *Synthesis* of standard discounted-Beta-Bernoulli mechanics with the schema’s environment-side parameters into a class-level structural cap. The constituent algebra is elementary — the discounted-update steady state $n_{\text{eff}} = 1/(1 - \lambda)$ is in Ljung 1987; the $\alpha = 1/(n + 1)$ form is in R Prop B.1. The AAT-distinctive contribution is the *combination* with ρ_{Σ}/R_{Σ} into a sharp reachability cap, surfaced via the §D.3 strengthen-first work that replaced the prior linear approximation (which had hidden the ceiling) with the exact form.

Related Work:

Table S.1

<i>ASF Concern</i>	<i>Prior-art Language</i>	<i>Relationship / Positioning</i>
Exponential forgetting / discounted least squares	Ljung 1987 <i>System Identification: Theory for the User</i> , MIT Press	<i>formal antecedent</i> — supplies the discounted-update mechanism that the C.1 steady state is built on
Class-level structural cap on gain-decay updates	NeurIPS 2026 Paper 2 §App-D ($\mathcal{A}_{\text{decay}}$ structural-class theorem)	<i>downstream sharpening</i> — lifts this Beta-Bernoulli instance to the gain-decay class; finite-gain mechanisms face bidirectional thresholds
Bias-variance trade-off in finite-gain adaptive filtering	Kushner & Yin 2003 <i>Stochastic Approximation and Recursive Algorithms and Applications</i> , 2nd ed. Springer	<i>contextual</i> — characterizes the upper bound complementing this lower bound for finite-gain (constant-step) mechanisms
Structural adaptation necessity (epistemic-layer analog)	4.8	<i>parallel structural pattern</i> — when parametric adaptation cannot keep up, structural change is required; this no-go is the strategic-layer instance

Appendix T

Derivation: Regret-Bound Derivation of the Strategy-Cost KL Direction

The variational form of the strategy-cost objective (9.9) carries a $D_{\text{KL}}(\pi^* \| Q_{\Sigma_t})$ relevance term. This appendix derives that **the π^* -first KL direction is forced** as the non-vacuous regret-bound form, and proves (conditional on a chain-rule additivity axiom motivated by 8.1) that **reverse-KL is uniquely forced within the direction-forced f-divergence family**.

Three sources of force compose. (i) The *direction* is forced decision-theoretically: bounded regret against an optimal policy is bounded above by value-range $\times$ total-variation, and by Pinsker’s inequality total variation is bounded by $\sqrt{D_{\text{KL}}(\pi^* \| Q)/2}$ in the π^* -first direction; the opposite-direction KL becomes infinite under deterministic optimal policy whenever the strategy places any mass off the optimum, giving a vacuous bound. (ii) Within the direction-forced family, the standard Hobson-Csiszár-Shore-Johnson uniqueness machinery via Aczél’s Cauchy functional equation forces reverse-KL *uniquely* among f-divergences once the chain-rule additivity axiom is granted — the divergence-layer instance of #disc-additive-coordinate-forcing’s meta-pattern (after chain-confidence decay; before log-odds at the update layer in #deriv-edge-update-natural-parameter). (iii) The bound is *sharpened* under AAT’s canonical deterministic-optimum scope: the **Bretagnolle-Huber identity** holds exactly (not as inequality), giving a two-sided regret bound $\Delta_{\min}(1 - e^{-D_{\text{KL}}}) \leq R \leq V_{\max}(1 - e^{-D_{\text{KL}}})$ strictly sharper than Pinsker — Lipschitz-equivalent to $(1 - e^{-D_{\text{KL}}})$ at both ends, with the Pinsker form becoming vacuous once $D_{\text{KL}} > 2$.

Secondary supporting properties of reverse-KL (gradient tractability, alignment with variational-inference machinery, MDL coding interpretation, Amari-Nagaoka Fisher geometry compatibility) are surfaced as *convergent evidence* rather than independent uniqueness grounds; the framework is explicit that Fisher-metric-at-second-order is not distinguishing within f-divergences (Eguchi 1983), so only the chain-rule axiom does the uniqueness work. The result is one of AAT’s cleanest demonstrations of the framework’s posture on coordinates: where a coordinate matters, the framework can point at the axiom that forces it; where the axiom isn’t AAT-internal, the framework can name what was adopted from outside.

Formal Expression.

§1 — Setup. Fix O_t , M_t , continuation policy π_{cont} , and horizon N_h . The action-value is the objective-functional evaluation at a : $V(a) := Q_O(M_t, a; \pi_{\text{cont}}, N_h)$ (6.6; V derives from O_t per 6.4). The optimal action is $a^* := \arg \max_a V(a)$; under AAT’s canonical scope (6.6), the optimal policy $\pi^*(\cdot | M_t) = \delta_{a^*}$ is the point mass on a^* (tied-optimum extensions §8). Write $V_{\max} := \max_a V(a) - \min_a V(a)$ for the value range at fixed M_t ; this is finite whenever $\mathcal{A}$ is a bounded action set under V .

$Q_{\Sigma_t}(\cdot \mid M_t)$ is the action distribution induced by the strategy DAG. In general Q_{Σ_t} is stochastic — the DAG’s edge-credence uncertainty induces policy stochasticity downstream of propagation.

§2 — Regret against the optimal policy.

The strategy-induced regret of Q_{Σ_t} against π^* is:

Definition (strategy-induced-regret).

$$R(Q_{\Sigma_t}) := V(a^*) - \mathbb{E}_{a \sim Q_{\Sigma_t}}[V(a)] = \sum_a Q_{\Sigma_t}(a) \cdot [V(a^*) - V(a)] \quad (\text{T.1})$$

Writing $\Delta(a) := V(a^*) - V(a) \in [0, V_{\max}]$ for the per-action regret gap. Three regret forms in the literature:

$$\mathbb{E}_{\pi^*}[V] - \mathbb{E}_{Q_{\Sigma_t}}[V], \quad V(a^*) - \mathbb{E}_{Q_{\Sigma_t}}[V], \quad \mathbb{E}_{\pi^*}[V - V_{Q_{\Sigma_t}}] \quad (\text{T.2})$$

coincide identically under deterministic π^* .

§3 — Total-variation bound (tight).

Derived (tv-regret-bound, from bounded V + deterministic π^).*

$$R(Q_{\Sigma_t}) = \sum_a Q_{\Sigma_t}(a) \Delta(a) \leq V_{\max} \cdot \sum_{a \neq a^*} Q_{\Sigma_t}(a) = V_{\max} \cdot (1 - Q_{\Sigma_t}(a^*)) \quad (\text{T.3})$$

For deterministic $\pi^* = \delta_{a^*}$:

$$\text{TV}(\pi^*, Q_{\Sigma_t}) = \frac{1}{2} \sum_a |\pi^*(a) - Q_{\Sigma_t}(a)| = 1 - Q_{\Sigma_t}(a^*) \quad (\text{T.4})$$

Therefore:

$$\boxed{R(Q_{\Sigma_t}) \leq V_{\max} \cdot \text{TV}(\pi^*, Q_{\Sigma_t})} \quad (\text{tight}) \quad (\text{T.5})$$

Tightness. Equality holds exactly when $\Delta(a) = V_{\max}$ for all $a \neq a^*$ (degenerate value landscape: every suboptimal action incurs the full value range). For typical landscapes the bound is loose by a factor $\mathbb{E}_{Q_{\Sigma_t}}[\Delta \mid a \neq a^*]/V_{\max} \in (0, 1]$.

§4 — Exact regret-reverse-KL identity under deterministic π^* . Under AAT’s canonical scope (deterministic $\pi^* = \delta_{a^*}$), reverse-KL and total variation are related by an *identity*, not merely an inequality. This gives a tight two-sided regret bound that the Pinsker inequality strictly weakens.

Identity. For deterministic $\pi^* = \delta_{a^*}$ and any Q with $Q(a^*) > 0$:

Derived (bb-identity-deterministic, under deterministic π^).*

$$\boxed{D_{\text{KL}}(\pi^* \| Q) = -\log Q(a^*) = -\log(1 - \text{TV}(\pi^*, Q))} \quad (\text{T.6})$$

Derivation. $D_{\text{KL}}(\delta_{a^*} \| Q) = \sum_a \delta_{a^*}(a) \log(\delta_{a^*}(a)/Q(a)) = -\log Q(a^*)$ (only the $a = a^*$ term contributes; $0 \log 0$ terms are conventionally 0). And $\text{TV}(\delta_{a^*}, Q) = 1 - Q(a^*)$ (direct from §3). Substituting gives the identity. $\square$

This is the Bretagnolle-Huber inequality (Bretagnolle & Huber 1978, “Estimation des densités,” *Séminaire de probabilités XII*, Springer LNM 649; Tsybakov 2009 §2.4; Sason & Verdú 2016) specialized to the deterministic- P case, where the general inequality $\text{TV} \leq \sqrt{1 - e^{-D_{\text{KL}}}}$ becomes an equality.

Combining with the tight TV-regret bound of §3:

Derived (exact-regret-reverse-kl, under deterministic π^).*

$$\boxed{R(Q_{\Sigma_t}) \leq V_{\max} \cdot (1 - e^{-D_{\text{KL}}(\pi^* \| Q_{\Sigma_t})})} \quad (\text{T.7})$$

Matching lower bound. Define the action-gap $\Delta_{\min} := \min_{a \neq a^*} \Delta(a) > 0$ (standard in RL theory; well-defined whenever the optimum is isolated over finite $\mathcal{A}$ with bounded V). Then:

$$R(Q_{\Sigma_t}) = \sum_{a \neq a^*} Q_{\Sigma_t}(a) \Delta(a) \geq \Delta_{\min} \cdot \sum_{a \neq a^*} Q_{\Sigma_t}(a) = \Delta_{\min} \cdot \text{TV}(\pi^*, Q_{\Sigma_t}) \quad (\text{T.8})$$

Under the exact BH identity:

$$R(Q_{\Sigma_t}) \geq \Delta_{\min} \cdot (1 - e^{-D_{\text{KL}}(\pi^* \| Q_{\Sigma_t})}) \quad (\text{T.9})$$

Lipschitz-equivalence corollary. Regret and $(1 - e^{-D_{\text{KL}}})$ are Lipschitz-equivalent with constants $\Delta_{\min}$ (below) and $V_{\max}$ (above):

Derived (matched-lower-bound, under deterministic π^ + isolated optimum).*

$$\frac{\Delta_{\min}}{V_{\max}} \leq \frac{R(Q_{\Sigma_t})}{V_{\max}(1 - e^{-D_{\text{KL}}})} \leq 1 \quad (\text{T.10})$$

The upper bound is tight when the value landscape is extremal ($\Delta(a) = V_{\max}$ for all $a \neq a^*$); the lower bound is tight when sub-optimal actions are uniformly bad ($\Delta_{\min} = \max_{a \neq a^*} \Delta(a)$).

$D_{\text{KL}}(\pi^* \| Q_{\Sigma_t}) \in [0, +\infty)$: finite and graded whenever $Q_{\Sigma_t}(a^*) > 0$; diverges only in the structural-failure limit where the strategy places zero mass on the optimum.

§4.1 — Pinsker bound as loose general form. The Pinsker inequality $\text{TV}(P, Q) \leq \sqrt{\frac{1}{2} D_{\text{KL}}(P \| Q)}$ (Tsybakov 2009 §2.4; Cover & Thomas 2006 §11.6) yields a regret bound that does not assume deterministic π^* :

$$R(Q_{\Sigma_t}) \leq V_{\max} \cdot \sqrt{\frac{1}{2} D_{\text{KL}}(\pi^* \| Q_{\Sigma_t})} \quad (\text{T.11})$$

Under AAT’s canonical scope, this is strictly weaker than the BH-identity form of §4. For any $D_{\text{KL}} > 0$: $1 - e^{-D_{\text{KL}}} < \sqrt{D_{\text{KL}}/2}$ for small D_{KL} (the BH form is linear in D_{KL} while Pinsker is $\sqrt{D_{\text{KL}}}$), and Pinsker’s $V_{\max} \sqrt{D_{\text{KL}}/2}$ exceeds the trivial $V_{\max}$ bound once $D_{\text{KL}} > 2$, giving vacuous content there. The BH identity is informative uniformly in D_{KL} .

Pinsker remains the correct tool for stochastic- π^* extensions (§9) where the deterministic- π^* scope is relaxed and the BH identity degrades back to inequality.

§5 — Direction-forcing claim (the load-bearing result).

Claim. The KL direction with π^* *first* is forced as the non-vacuous regret-bound form. Forward-KL $D_{\text{KL}}(Q_{\Sigma_t} \| \pi^*)$ is *not* a non-vacuous bound on $R(Q_{\Sigma_t})$ under deterministic π^* .

Derived (kl-direction-forced, from deterministic π^).*

Derivation. Expanding:

$$D_{\text{KL}}(Q_{\Sigma_t} \| \pi^*) = \sum_a Q_{\Sigma_t}(a) \log(Q_{\Sigma_t}(a) / \pi^*(a)) \quad (\text{T.12})$$

For any $a \neq a^*$: $\pi^*(a) = 0$ (point mass assumption). The summand $Q_{\Sigma_t}(a) \log(Q_{\Sigma_t}(a)/0) = +\infty$ unless $Q_{\Sigma_t}(a) = 0$. Therefore $D_{\text{KL}}(Q_{\Sigma_t} \| \pi^*) = +\infty$ whenever Q_{Σ_t} places any mass off a^* — for *all but a measure-zero subset* of stochastic strategies. A bound “ $R \leq +\infty$ ” is vacuous.

This is the same shape of degeneracy as the original Shannon-MI form (Gemini Finding 2) that the V-medium move was introduced to escape: same degeneracy-when- π^* -is-deterministic, different value (0 vs $+\infty$). Only the *reverse* direction escapes both.

Alignment with the structural problem. Forward-KL is natural for *mode-covering* inference (where Q is asked to cover P ’s full support) and variational inference when Q is learned to match a full distribution. Reverse-KL is natural for *mode-seeking* (concentrate Q on P ’s mode). The AAT strategy problem is mode-seeking by construction: the strategy should concentrate on a^* . The direction alignment with the structural problem is not accidental.

Asymmetry is forced by regret’s one-sidedness. A second, independent argument for asymmetry — one that does not rely on the chain-rule axiom of §6.1. Under deterministic π^* , regret is a *one-sided* quantity: $R(Q_{\Sigma_t}) = \sum_{a \neq a^*} Q_{\Sigma_t}(a) \Delta(a)$ contributes only from Q_{Σ_t} ’s off-optimum mass; π^* ’s “mass off Q_{Σ_t} ” is vacuous since π^* has no support off a^* . Any divergence whose role is to bound this quantity should therefore be asymmetric — it should penalize Q_{Σ_t} ’s deviation from π^* without symmetrically penalizing the (trivially zero) deviation of π^* from Q_{Σ_t} . Symmetric divergences (squared Hellinger, Jensen-Shannon, symmetrized KL) treat the two deviations interchangeably; under the one-sided regret quantity this is operationally wrong — it introduces a term with no semantic role. The asymmetry requirement is thus *structural*, emerging directly from what regret is as a functional, and is established independently of the chain-rule axiom that picks the specific asymmetric form in §6.1. The two arguments compose: *direction-forcing* (vacuity under deterministic π^*) + *asymmetry-forcing* (regret is one-sided) $\Rightarrow$ the admissible bounding divergence is π^* -first and asymmetric; chain-rule additivity (§6.1) then picks reverse-KL uniquely within that family.

Terminology note (Bishop-vs-AAT “reverse-KL”). AAT’s canonical reverse-KL is $D_{\text{KL}}(\pi^* \| Q_{\Sigma_t})$ — optimum-first, agent-second. Under Bishop 2006’s convention ($D_{\text{KL}}(q \| p)$ with q the approximation is called “reverse-KL”), AAT’s form would be labeled “forward-KL.” The two naming conventions disagree in surface labeling but agree in the operational property — both pick out the *mode-seeking* direction (concentrate the approximation on the target’s mode).

Literature that uses the Bishop convention (variational inference, generative modeling, Levine 2018) and literature that uses the π^* -first / target-first convention (decision theory, regret analysis, Rubin-Shamir-Tishby 2012) describe the same operational quantity under deterministic π^* . When this segment writes “reverse-KL,” the π^* -first / AAT convention is meant; under softened targets the two conventions describe different objects and must be disambiguated explicitly.

§6 — Admissible-divergence family and uniqueness of reverse-KL.

The regret-bound argument forces the *direction* (the reference distribution π^* appears first) but admits a family of divergences. Each member yields a valid regret bound; they differ in tightness and operational properties:

Discussion (admissible-regret-divergences).

Table T.1

<i>Divergence</i>	<i>Bound on R</i>	<i>Tightness</i>	<i>Finite under det. π^*?</i>	<i>Gradient-tractable?</i>
$\text{TV}(\pi^*, Q_{\Sigma_t})$	$V_{\max} \cdot \text{TV}$	Tight (extremal V)	Yes	No (non-differentiable)
$D_{\text{KL}}(\pi^* \ Q_{\Sigma_t})$ via Pinsker	$V_{\max} \sqrt{\frac{1}{2} D_{\text{KL}}}$	Loose by $\sqrt{\cdot}$	Yes	Yes
$D_{\text{KL}}(\pi^* \ Q_{\Sigma_t})$ via Bretagnolle-Huber	$V_{\max} \sqrt{1 - e^{-D_{\text{KL}}}}$	Tighter than Pinsker for large D_{KL}	Yes	Yes
$\chi^2(\pi^* \ Q_{\Sigma_t})$ (Le Cam)	$V_{\max} \cdot \frac{1}{2} \sqrt{\chi^2}$	Typically looser than Pinsker	Yes: $\chi^2 =$ $1/Q_{\Sigma_t}(a^*) - 1$	Yes
$D_{\alpha}(\pi^* \ Q_{\Sigma_t})$ (Rényi, $\alpha \geq 1$)	Various	Interpolates KL ($\alpha \rightarrow 1$) and χ^2 ($\alpha = 2$)	Yes for $\alpha \geq 1$	Yes
$D_{\text{KL}}(Q_{\Sigma_t} \ \pi^*)$ (forward-KL)	$+\infty$	Vacuous	No	—

What is uniquely forced by the regret-bound argument alone. The direction: the reference π^* is first. This is a real derivation outcome, not a selection. Within the π^* -first family, multiple f-divergences each give valid bounds; an additional structural axiom is required to pick one uniquely.

§6.1 — Chain-rule uniqueness theorem.

Theorem (chain-rule / additivity uniqueness of KL among f-divergences; folk theorem, standard functional-equation derivation). Let $D_f(P \| Q) = \sum_x Q(x) f(P(x)/Q(x))$ be a smooth f-divergence with f convex and $f(1) = 0$. The chain rule

Derived (reverse-kl-uniqueness, Conditional on chain-rule axiom).

$$D_f(R_{XY} \| Q_{XY}) = D_f(R_X \| Q_X) + \mathbb{E}_{P_X}[D_f(R_{Y|X} \| Q_{Y|X})] \quad (\text{T.13})$$

holds for all joint distributions (X, Y) if and only if $f(t) = c \cdot t \log t$ for some $c > 0$ — i.e., D_f is reverse-KL up to positive scaling.

References. The theorem is a classical folk result obtainable by direct functional-equation argument (sketched below). The canonical published axiomatic characterizations equivalent to the chain-rule statement are: **Hobson 1969**, “A new theorem of information theory,” *J. Stat. Phys.* 1(3):383–391 — uniqueness of the Kullback expression via a composition/additivity axiom; **Csiszár 1991**, “Why least squares and maximum entropy? An axiomatic approach to inference

for linear inverse problems,” *Annals of Statistics* 19(4):2032–2066 — Theorem 3 corollary: the only transitive statistical projection rule is the I-divergence projection rule; Theorem 5: product-consistency characterizes I-divergence uniquely. See also **Shore & Johnson 1980**, “Axiomatic derivation of the principle of maximum entropy and the principle of minimum cross-entropy,” *IEEE Trans. Info. Theory* 26(1):26–37 (system-independence axiom); **Sanov 1957**, “On the probability of large deviations of random variables,” *Mat. Sb.* 42(84):11–44 (large-deviation rate function for empirical-distribution concentration); **Aczél & Daróczy 1975**, *On Measures of Information and Their Characterizations* (Academic Press) for the general functional-equation machinery.

These are not independent uniqueness routes. Hobson’s composition axiom, Csiszár’s transitive-projection axiom, Shore-Johnson’s system-independence axiom, and Sanov’s sampling-consistency condition are *structurally-equivalent reformulations* of the same underlying requirement — all factor through independence-of-sub-problems (joint inference on independent subsystems decomposes additively over factorizations). The Cauchy functional equation that each reduces to is the common structural content beneath varied surface formulations. This matters for #disc-additive-coordinate-forcing: the chain-rule axiom’s load-bearing role is not weakened by noting multiple canonical references; conversely, the multiple references do not provide multiple independent uniqueness arguments. No known uniqueness route outside the independence-on-sub-problems family exists.

Proof sketch. Write $r_x = P(x)/Q(x)$, $s_{y|x} = P(y | x)/Q(y | x)$. The chain-rule identity must hold for all joint decompositions, which forces the functional equation $f(rs) = f(r) + rf(s) + g(r)$ for all $r, s > 0$. Combined with $f(1) = 0$ and convexity, the unique solution is $f(t) = c \cdot t \log t$ for $c > 0$ (Aczél & Daróczy 1975 §4).

Why other members of the family fail the chain rule. Concrete counterexample for χ^2 : take Q_X uniform on $\{x_1, x_2\}$, $P_X = (3/4, 1/4)$, $Q(y | x)$ uniform, $P(y | x) = (3/4, 1/4)$ for both x . Direct calculation gives $\chi^2(P_{XY} \| Q_{XY}) = 9/16$, while $\chi^2(P_X \| Q_X) + \mathbb{E}_{P_X}[\chi^2(P_{Y|X} \| Q_{Y|X})] = 1/4 + 1/4 = 8/16$. Non-additive: $9/16 \neq 8/16$. Rényi- α for $\alpha \neq 1$ fails analogously; Bretagnolle-Huber is a monotone transform of reverse-KL (not an independent f-divergence); Hellinger-squared ($\alpha = 0$) likewise fails.

AAT-internal motivation for the chain-rule axiom. The chain rule is the *divergence-level analog* of the additive log-confidence decay in #der-chain-confidence-decay. AAT has already committed to additive-mismatch-decomposition-along-causal-chains in that segment; the divergence-level parallel is:

$$D(\pi^* \| Q_{\Sigma_t}) = \sum_{t=1}^T \mathbb{E}_{\pi^*} \left[D(\pi^*(\cdot | a_{<t}^*) \| Q(\cdot | a_{<t}^*)) \right] \quad (\text{T.14})$$

— total mismatch between π^* and Q_{Σ_t} decomposes additively over the DAG’s causal layers along the optimal trajectory. Non-chain-rule divergences (e.g., χ^2) give super-additive decompositions in which layer-mismatches amplify multiplicatively — structurally discordant with the DAG factorization (8.3). Adopting the chain-rule axiom is therefore not arbitrary; it is the divergence-level version of a decomposition principle AAT already relies on.

Corollary (within the chain-rule axiom). Under deterministic π^* + bounded value + chain-rule additivity, reverse-KL is the unique smooth f-divergence in the direction-forced admissible family. The “Reverse-KL canonical among smooth divergences” status upgrades from formulation choice (under the pre-axiom reading) to *Derived (conditional on the chain-rule axiom)*.

What the uniqueness theorem does not do. It fixes the divergence within the f-divergence family; it does not fix the bounding function g around the divergence (Pinsker vs Bretagnolle-Huber vs Le Cam-on- χ^2 are *different bounds using different divergences*, so only Pinsker and BH-on-KL remain in scope after uniqueness — and these differ only in the g -envelope on top of reverse-KL, not in the divergence itself). It does not resolve the linear-vs-square-root form question of §7 (that is a Lagrangian-shape choice orthogonal to the divergence choice). Monotone transforms of reverse-KL (e.g., $1 - e^{-D_{\text{KL}}}$) are equivalence-class members for gradient-based optimization and are not ruled out.

Scope of the axiom. The chain rule is stated for arbitrary joint decompositions. AAT’s strategy spaces are discrete at the proposition level; the theorem applies directly. For continuous strategy spaces (e.g., low-level control), the chain rule

extends by replacing sums with integrals against the dominating measure; the functional-equation derivation is identical in the continuous case (see Liese & Vajda 1987, *Convex Statistical Distances*, Teubner, for the standard measure-theoretic treatment of f-divergences and their decomposition properties).

§6.2 — *Secondary characterizations (now supporting rather than load-bearing)*. Four additional properties further motivate reverse-KL as the canonical choice; under §6.1 they are no longer load-bearing for uniqueness but remain informative about why reverse-KL is a comfortable fit with the rest of the AAT framework:

1. **Gradient-tractability.** TV is non-differentiable at mass boundaries; reverse-KL is smooth. Operational minimization via gradient methods rules out TV despite its tightness — independent of the uniqueness argument.
2. **Variational-inference alignment.** Reverse-KL is the standard divergence in the variational-inference lineage (ELBO derivation uses reverse-KL; Friston et al. 2017; Da Costa et al. 2020). The V-medium move’s rhetorical payoff — shared vocabulary with active inference — lives on reverse-KL specifically. This is *convergent evidence* that reverse-KL is natural in multiple frames, not an independent uniqueness argument: the ELBO decomposition itself uses the chain rule as a sub-step, so the logical priority runs chain rule $\rightarrow$ reverse-KL $\rightarrow$ ELBO.
3. **Fisher geometry.** Reverse-KL’s second-order expansion gives the Fisher information metric (Amari & Nagaoka 2000). *But by the differential-geometric framework in Eguchi 1983* (“Second order efficiency of minimum contrast estimators in a curved exponential family,” *Annals of Statistics* 11(3):793–803, §2 contrast-function development; modern restatement at Amari & Cichocki 2010 Theorem 5 eq. (126)), every smooth f-divergence with $f''(1) > 0$ induces the Fisher metric at second order up to the scalar $f''(1)$. Fisher-metric-at-second-order is *not* a distinguishing axiom within the f-divergence family — it is satisfied by every member. Noted here because it is a commonly-invoked distinguishing property; it does not distinguish. (Eguchi’s Theorem 3 itself is about estimator efficiency via Γ^1 -transversality; the f-divergence/Fisher-metric result is a consequence of the §2 contrast-function machinery that supports Theorem 3, not the theorem statement.)
4. **Information-budget / MDL interpretation.** $D_{KL}(\pi^* \| Q_{\Sigma_t}) = -\log Q_{\Sigma_t}(a^*)$ is the expected extra bits needed to code samples from π^* under Q_{Σ_t} ’s code. The coding interpretation aligns with the segment’s MDL framing. MDL is itself a chain-rule-respecting coding scheme (the bits-to-code-a-joint decompose additively over factorizations); consistency with MDL is convergent with §6.1.

Čencov/Morozova-Chentsov background (Čencov 1982 *Statistical Decision Rules and Optimal Inference*, AMS Translations of Mathematical Monographs 53; Morozova & Chentsov 1991, “Markov invariant geometry on state manifolds,” *Itogi Nauki i Tekhniki*, translated in *J. Sov. Math.* 56(5):2648–2669; Ay, Jost, Lê & Schwachhöfer 2017, *Information Geometry*, Springer). The f-divergence family itself is characterized by Markov-morphism invariance (sufficient-statistic coarse-graining invariance) — this is what makes “f-divergences” the right background family to search within. Amari 2009 shows the alpha-divergences are the sub-family that are simultaneously f-divergences and Bregman divergences. Reverse-KL sits at $\alpha = 1$ within the alpha-family; the chain-rule axiom of §6.1 is what picks $\alpha = 1$ out of the one-parameter alpha-family.

§6.3 — *Bregman-Fenchel identification: reverse-KL and log-odds as dual coordinates.*

On the probability simplex Δ^{n-1} , the **negative-entropy potential** $\phi(Q) = \sum_a Q_a \log Q_a$ is strictly convex, Legendre, and essentially smooth on the relative interior (Rockafellar 1970 *Convex Analysis* §26). Its Fenchel conjugate is the **log-partition function** $\phi^*(\eta) = \log \sum_a e^{\eta_a}$ on the natural-parameter space $\mathbb{R}^n$ modulo the affine-gauge direction. The primal-dual correspondence is softmax: $Q = \nabla \phi^*(\eta)$ gives $Q_a = e^{\eta_a} / \sum_{a'} e^{\eta_{a'}}$, and $\eta = \nabla \phi(Q) = \log Q + \mathbf{1}$ (up to the constraint-normal direction) so $\eta_a - \eta_b = \log(Q_a/Q_b)$ is the **log-odds ratio**.

Derived (bregman-fenchel-dual-pair, exact; standard Legendre-Fenchel).

The Bregman divergence induced by ϕ on the primal simplex is **reverse-KL** exactly:

$$B_\phi(P, Q) := \phi(P) - \phi(Q) - \langle \nabla \phi(Q), P - Q \rangle = \sum_a P_a \log \frac{P_a}{Q_a} = D_{\text{KL}}(P \| Q) \quad (\text{T.15})$$

Derivation. Expand: $\phi(P) - \phi(Q) = \sum_a (P_a \log P_a - Q_a \log Q_a)$; the inner-product term using $\nabla \phi(Q)_a = \log Q_a + 1$ gives $\sum_a (\log Q_a + 1)(P_a - Q_a) = \sum_a (\log Q_a)(P_a - Q_a)$ since $\sum_a (P_a - Q_a) = 0$. Substituting and simplifying yields $\sum_a P_a \log(P_a/Q_a) = D_{\text{KL}}(P \| Q)$. $\square$

Identification with #deriv-edge-update-natural-parameter. The log-odds coordinate derived under evidential-additivity in #deriv-edge-update-natural-parameter is the *Fenchel-dual natural parameter* of this Legendre-Fenchel pair. The divergence-layer coordinate of this segment (reverse-KL on Δ^{n-1}) and the update-layer coordinate of #deriv-edge-update-natural-parameter (log-odds on $\mathbb{R}^n/\mathbf{1}$) are two sides of one exponential-family geometric structure, viewed through primal and dual coordinates respectively.

What this adds. The two segments independently derive their coordinates via Cauchy-FE on logically-independent AAT-internal axioms (chain-rule additivity here; evidential additivity in #deriv-edge-update-natural-parameter). The Fenchel-Bregman correspondence shows that *the coordinates the two axioms pick out are related by Legendre duality* — the axioms are not redundant (they have different logical forms and constrain different objects; see §6.1 vs. #deriv-edge-update-natural-parameter’s axiom) but the forced coordinates *coincide on one geometric object*. This is a Discussion-grade observation about the relationship between two primary instances of #disc-additive-coordinate-forcing; it is textbook Legendre-Fenchel (Amari & Nagaoka 2000 §3.5; Bregman 1967; Beck & Teboulle 2003 *Mirror descent and nonlinear projected subgradient methods* — relevant for the operational consequence that mirror-descent on the policy simplex with negative-entropy regularizer is equivalent to exponentiated-gradient descent on log-odds coordinates). The full geometric-unification reading sits in #disc-additive-coordinate-forcing’s meta-treatment.

§6.4 — Information-theoretic-MDP lineage and AAT’s direction choice.

AAT’s reverse-KL strategy-cost form sits within an established research lineage — the information-theoretic treatment of decision problems initiated by **Tishby & Polani 2011** (“The Information Theory of Decision and Action,” in *Perception-Action Cycle: Models, Architectures, and Hardware*, Springer, pp. 601–636; Eq. 15 p. 19 introduces **Information-to-Go** $\mathfrak{I}^\pi(s_t, a_t)$, the conditional multi-information of the entire future state-action trajectory; Eq. 17 p. 21 the trade-off objective $\arg \min_\pi [\mathfrak{I}^\pi - \beta Q^\pi]$ and developed by **Rubin, Shamir & Tishby 2012** (“Trading value and information in MDPs,” in *Decision Making with Imperfect Decision Makers*, Springer, pp. 57–74; §2.2 Eq. 3 p. 4 **control information** $\Delta I(s) = \sum_a \pi_s(a) \log(\pi_s(a)/\rho_s(a))$; §3.1 Theorem 1 p. 5 Bellman recursion for the free-energy $F_\pi = I_\pi - \beta V_\pi$). The control-as-inference lineage via **Levine 2018** (“Reinforcement Learning and Control as Probabilistic Inference,” arXiv:1805.00909; §3.1 Eq. 11 and §5.2 p. 15) develops the max-entropy-RL framework using $D_{\text{KL}}(p_\theta \| p_{\text{tgt}})$ throughout. The shared framework — value and information as joint first-class quantities in MDP optimization, with a KL divergence measuring the information cost of a policy against a reference — is the context in which AAT’s strategy-cost objective naturally lives.

Discussion (information-theoretic-mdp-lineage), positioning rather than derivation.

AAT’s direction is distinctive within this lineage. Tishby-Polani 2011 uses the stationary marginal $\pi(a)$ of the policy itself as reference (Eq. 16 p. 20, Eq. 22 p. 22) — a direction-neutral multi-information rather than a KL-to-external-reference. Rubin et al. 2012 uses a default policy $\rho_s(a)$ (typically uniform) with direction $D_{\text{KL}}(\pi_s \| \rho_s)$ — **agent-first, default-second**, a regularization/complexity cost. Levine 2018 uses $D_{\text{KL}}(p_\theta \| p_{\text{tgt}})$ — **proposal-first, target-second**, a variational ELBO form with softened exponentiated-reward target. **AAT’s** $D_{\text{KL}}(\pi^* \| Q_{\Sigma_t})$ — **optimum-first, agent-second** — is the *opposite* direction to all three.

The direction is forced, not chosen. AAT’s direction is not a convention adopted from the surrounding literature; it is forced by two independent structural requirements of this segment’s derivation: (i) the regret-bound argument of §5 requires the optimum as the reference for the bound to be tight (mode-seeking toward a^*); (ii) the Bretagnolle-Huber identity of §4 is *exact* in the π^* -first direction under deterministic π^* and degenerate in the flipped direction. The

flipped direction (Rubin 2012's / Levine 2018's form with $\rho = \pi^*$) is vacuous here for the same reason forward-KL is vacuous in §5: π^* has no support off a^* . AAT *owns* the direction-distinctiveness rather than inheriting it.

Rubin 2012 Theorem 3 PAC-Bayesian motivation. A fourth independent operational property of the π^* -first KL form comes from Rubin-Shamir-Tishby 2012 Theorem 3 (p. 10): under a-priori stochastic default policy and empirical reward estimate, the control information I_π bounds the generalization gap of the value function:

$$\tilde{D}_{\text{KL}}(\hat{V}_\pi(s_0) \parallel V_\pi(s_0)) \leq \frac{I_\pi(s_0) + \log(2m/\delta)}{m-1} \quad (\text{T.16})$$

This is a PAC-Bayesian concentration bound. In AAT terms: policies with lower strategy-cost have tighter value-function generalization under finite empirical data. Together with (a) the regret-bound direction-forcing of §5, (b) the chain-rule uniqueness of §6.1, and (c) the Bregman-Fenchel dual identification of §6.3, this is a fourth independent motivation for the KL-to-reference form — operational rather than axiomatic, and independent of the chain-rule axiom. The direction differs (Rubin uses agent-first; AAT uses optimum-first), but the generalization-bound structure transfers by the same PAC-Bayesian machinery: I_π as regularizer against a reference.

Relation to #form-information-bottleneck. The canonical IB objective is $I(X; T) - \beta I(T; Y)$ with Shannon mutual information on both terms. AAT's strategy-cost relevance term is $D_{\text{KL}}(\pi^* \parallel Q_{\Sigma_t})$ — a KL divergence to a target, not a mutual information. The two are *sibling lineages*, both descended from Shannon rate-distortion theory (the IB instance uses MI-to-relevant-label as the fidelity term; the information-theoretic-MDP instance uses KL-to-reference-policy); neither reduces to the other without a change of relevance variable. AAT's compression-operations framework (#disc-compression-operations) uses the IB form for M_t , G_t^{shared} , and Λ compressions; the strategy-cost compression uses the information-theoretic-MDP form. Both are valid Lagrangian relaxations of an underlying rate-distortion problem; the choice of fidelity term depends on whether the compressed variable is meant to preserve information about an observable (IB) or to match a reference policy (information-theoretic-MDP). The “abandoned IB purity” concern raised in external reviews (Gemini 2026-04-24) is defused at this level: AAT has not abandoned IB; its strategy-cost compression uses a different relevance quantity than its model-compression, and both are principled under the rate-distortion umbrella.

§7 — β_Σ interpretation under the regret bound.

The Pinsker bound gives $R(Q_{\Sigma_t}) \leq V_{\max} \sqrt{\frac{1}{2} D_{\text{KL}}(\pi^* \parallel Q_{\Sigma_t})}$. Two scale forms for the segment's trade-off parameter β_Σ arise:

Discussion (β_Σ -local-vs-global).

(a) **Square-root-KL form** (tight regret-bound scale):

$$\mathcal{L}(\Sigma_t) = \mathcal{R}(\Sigma_t) + \beta_\Sigma \cdot \sqrt{D_{\text{KL}}(\pi^* \parallel Q_{\Sigma_t})}, \quad \beta_\Sigma = V_{\max}/\sqrt{2} \quad (\text{T.17})$$

β_Σ is *globally naturalized* as a constant scale proportional to the value range at fixed M_t . Cost: the form departs from the linear-Lagrangian IB shape that AA uses across the four compression operations.

(b) **Linear-KL form** (IB-shape instance, preserved in the segment):

$$\mathcal{L}(\Sigma_t) = \mathcal{R}(\Sigma_t) + \beta_\Sigma \cdot D_{\text{KL}}(\pi^* \parallel Q_{\Sigma_t}) \quad (\text{T.18})$$

Under this form, β_Σ 's regret-bound interpretation is *local*: differentiating the Pinsker bound, $\partial R / \partial D_{\text{KL}} = V_{\max} / (2\sqrt{2D_{\text{KL}}})$, so β_Σ represents the local cost-per-bit at the operating point but varies with D_{KL} . Upside: consistency with the IB linear-Lagrangian shape (AA). Downside: no uniform global β_Σ scale.

Choice made in 9.9. Keep the linear form (IB-shape alignment); note the square-root form in the Epistemic Status as the “source” derivation. The direction-forcing claim is the load-bearing strengthening; the linear-vs-square-root choice is a second-order trade-off that preserves architectural consistency at the cost of a fully naturalized β_Σ .

§8 — Bound tightness and scope limits. Vacuity regimes (where the regret bound fails to be informative):

1. $V_{\max} = \infty$ (**unbounded value**). Then all bounds are ∞ . AAT’s O_t is $\mathbb{R}$ -valued (6.4); boundedness of V over $\mathcal{A}$ at fixed M_t is an additional assumption. Stated here as a scope condition of the derivation.
2. $Q_{\Sigma_t}(a^*) = 0$ (**strategy cannot express the optimum**). Then $D_{\text{KL}}(\pi^* \| Q_{\Sigma_t}) = +\infty$. This is structural failure — the strategy’s DAG cannot produce the optimal action — and the infinite KL correctly flags it. *Informative degeneracy*, not pathology: operational minimization rejects such Σ_t .
3. π^* **not deterministic (outside canonical scope)**. For stochastic π^* with larger support than Q_{Σ_t} , reverse-KL may still be infinite. §9 addresses tied-optimum extensions; stochastic π^* under softmax-smoothing is future work.

Tightness comparison. Under deterministic π^* (AAT’s canonical scope), the Bretagnolle-Huber **identity** $D_{\text{KL}}(\pi^* \| Q_{\Sigma_t}) = -\log(1 - \text{TV}(\pi^*, Q_{\Sigma_t}))$ makes the KL and TV bounds *operationally equivalent in reverse-KL coordinates*: the KL-based bound $R \leq V_{\max}(1 - e^{-D_{\text{KL}}})$ of §4 is exactly the TV-based bound of §3 re-expressed in the KL coordinate, with matching lower bound $R \geq \Delta_{\min}(1 - e^{-D_{\text{KL}}})$ (§4). Pinsker is strictly weaker under this scope (see §4.1). Outside the deterministic- π^* scope, the BH identity degrades back to the inequality $\text{TV} \leq \sqrt{1 - e^{-D_{\text{KL}}}}$ (Bretagnolle-Huber 1978; tighter than Pinsker for large D_{KL} but no longer an identity); Pinsker is the textbook fallback for general distributions. The net: reverse-KL is not merely “the tightest smooth divergence” — in AAT’s canonical scope it carries the *exact* regret-bound relationship via BH, with Pinsker as a loose alternative retained only for stochastic- π^* extensions.

§9 — Extensions. Tied-optimum case. If π^* has support on a tied-optimum set $\mathcal{A}^* = \{a : V(a) = V(a^*)\}$ with uniform mass, reverse-KL is finite whenever Q_{Σ_t} covers $\mathcal{A}^*$. The regret-bound argument extends directly: $R(Q_{\Sigma_t}) = \sum_a Q_{\Sigma_t}(a) \Delta(a) \leq V_{\max} \cdot \mathbb{P}_{Q_{\Sigma_t}}(a \notin \mathcal{A}^*)$ and Pinsker applies unchanged.

Softmax-smoothed π^* (stochastic π^* for non-degeneracy reasons). A regret bound of the form

$$\mathbb{E}_{\pi^*}[V] - \mathbb{E}_{Q_{\Sigma_t}}[V] \quad (\text{T.19})$$

with softmax-weighted π^* also admits the Pinsker-KL step, with reverse-KL again the non-vacuous direction (forward-KL remains vacuous whenever π^* has wider support than Q_{Σ_t} , and vice versa). Deferred for explicit treatment.

What Is Derived vs. What Is Chosen.

Table T.2

Property	Source	Strength
Regret definition $R(Q_{\Sigma_t}) := V(a^*) - \mathbb{E}_{Q_{\Sigma_t}}[V]$	Definitional; collapses to the three literature forms under deterministic π^*	Definition
TV-regret bound $R \leq V_{\max} \cdot \text{TV}(\pi^*, Q_{\Sigma_t})$	Bounded value range + deterministic π^*	Proved (tight)
BH identity $D_{\text{KL}}(\pi^* \ Q_{\Sigma_t}) = -\log(1 - \text{TV}(\pi^*, Q_{\Sigma_t}))$ under deterministic π^*	Direct calculation (Bretagnolle-Huber 1978 specialized to $P = \delta_{a^*}$)	Proved (exact identity)
Regret-reverse-KL bound $R \leq V_{\max}(1 - e^{-D_{\text{KL}}(\pi^* \ Q_{\Sigma_t})})$	TV-regret bound + BH identity, under deterministic π^*	Proved (tight in upper direction)

<i>Property</i>	<i>Source</i>	<i>Strength</i>
Matching lower bound $R \geq \Delta_{\min}(1 - e^{-D_{\text{KL}}(\pi^* \ Q_{\Sigma_t})})$ Pinsker-KL bound $R \leq V_{\max} \sqrt{\frac{1}{2} D_{\text{KL}}(\pi^* \ Q_{\Sigma_t})}$	Isolated-optimum action-gap $\Delta_{\min} > 0$ + BH identity Pinsker's inequality applied to TV bound	Proved (Lipschitz-equivalence with $V_{\max} / \Delta_{\min}$ constants) Proved (strictly weaker than BH identity under deterministic π^* ; correct form for stochastic- π^* extensions)
KL direction forced (reverse-KL, π^* -first)	Forward-KL is vacuous ($+\infty$) under deterministic π^* whenever Q_{Σ_t} has off-optimum mass	Proved (direction uniquely forced)
Asymmetry forced by regret's one-sidedness	Regret contributes only from Q_{Σ_t} 's off-optimum mass; π^* 's off- Q deviation is vacuous; symmetric divergences penalize interchangeably, which is operationally wrong	Proved (asymmetry forced; independent of chain-rule axiom)
Reverse-KL uniquely forced within direction-forced f-divergences	Chain-rule additivity axiom (Hobson 1969; Csiszár 1991 Theorem 3 corollary and Theorem 5; standard functional-equation derivation per Aczél & Daróczy 1975); AAT-internally motivated as divergence-level analog of 8.1 (§6.1)	Derived (conditional on chain-rule axiom)
Secondary characterizations (gradient-tractability, VI-alignment, MDL coding)	Independent operational properties; each compatible with reverse-KL, none individually distinguishing	Formulation support (non-load-bearing under §6.1)
Fisher-metric-at-second-order does not distinguish reverse-KL	Eguchi 1983 §2 contrast-function framework (<i>Ann. Statist.</i> 11(3):793–803): every smooth f-divergence with $f''(1) > 0$ induces the Fisher metric up to scalar $f''(1)$	Proved (no-go for Path A as uniqueness axiom)
Admissible family members (TV, Bretagnolle-Huber, χ^2 , Rényi- α)	Each yields a valid regret bound; survey §6	Derived (each)
Under deterministic π^* : TV and KL-coordinate bounds are equivalent via BH identity; Pinsker strictly weaker	§4 BH identity + §4.1 Pinsker comparison	Proved (scope-dependent tightness ordering)
Reverse-KL is the Bregman divergence of negative-entropy on Δ^{n-1} ; log-odds is the Fenchel-dual natural coordinate (Legendre-Fenchel pair)	Standard Legendre-Fenchel (Rockafellar 1970 §26; Amari & Nagaoka 2000 §3.5; Bregman 1967) applied to $\phi(Q) = \sum_a Q_a \log Q_a$	Proved (textbook identification; §6.3)
AAT's π^* -first KL direction is distinctive within the information-theoretic-MDP lineage (TP2011 / Rubin 2012 / Levine 2018 all put agent-first) and is <i>forced</i> (not inherited) by the regret-bound + BH-identity derivations	Literature positioning (§6.4) + derivation structure (§§4–5)	Discussion (positioning; direction forced by §§4–5 derivations)
PAC-Bayesian generalization bound: $\tilde{D}_{\text{KL}}(V_{\pi} \ V_{\pi}) \leq (I_{\pi} + \log(2m/\delta))/(m-1)$	Rubin-Shamir-Tishby 2012 Theorem 3	Derived (external theorem applied; independent operational motivation alongside regret-bound + chain-rule + Fenchel-dual)
$\beta_{\Sigma} \propto V_{\max}$ naturalization (square-root-KL form)	Direct identification from Pinsker bound	Derived (under square-root form)
β_{Σ} local interpretation (linear-KL form, segment-retained)	Differentiation of Pinsker bound at operating point	Derived (local only)
Linear-KL form retained over square-root form	IB-shape alignment with AA	Formulation choice
Vacuity regimes ($V_{\max} = \infty$, $Q_{\Sigma_t}(a^*) = 0$, stochastic π^*)	Direct analysis of the bound	Proved (boundary)
Tied-optimum extension	π^* uniform on $\mathcal{A}^*$; bound adapts directly	Derived
Softmax-smoothed π^* extension	Sketched in §9	Hypothesis (deferred)

Property	Source	Strength
Uniqueness under chain-rule axiom (§6.1)	Hobson 1969 / Csiszár 1991 Theorem 3 corollary and Theorem 5 / standard functional-equation argument per Aczél & Daróczy 1975, applied to the direction-forced family; axiom AAT-internally motivated (§6.1)	Derived (conditional)

The dividing line: the KL **direction** is forced by the regret-bound derivation (strong result). The specific reverse-KL **form** within the direction-forced f-divergence family is *derived under the chain-rule additivity axiom* (§6.1) — conditional on an axiom that AAT independently motivates as the divergence-level version of the additive-decomposition principle already in 8.1. The β_Σ naturalization is partial — available globally only under the square-root form, locally under the linear form retained for IB-shape alignment.

Epistemic Status.

Robust qualitative overall, with several specific claims at *exact* tier. The direction-forcing claim (§5) is derived under standard assumptions (bounded value, deterministic π^* on the canonical AAT scope). The **Bretagnolle-Huber identity under deterministic π^*** (§4, $D_{\text{KL}}(\pi^* \| Q) = -\log(1 - \text{TV})$) is *exact* by direct calculation, and yields the tight regret bound $R \leq V_{\max}(1 - e^{-D_{\text{KL}}})$ with matching lower bound $\Delta_{\min}(1 - e^{-D_{\text{KL}}})$ — the segment’s primary bound form. Pinsker (§4.1) is retained as the strictly-weaker general form, correct for stochastic- π^* extensions. The TV bound (§3) is textbook. The admissible-divergence family characterization (§6) is a truthful survey. The **chain-rule uniqueness theorem** (§6.1) upgrades the reverse-KL selection to “uniquely forced conditional on the chain-rule axiom”; the axiom is AAT-internally motivated as the divergence-level analog of #der-chain-confidence-decay. **Asymmetry-from-one-sidedness** (§5) provides a second independent argument for the direction-forcing side, not requiring the chain-rule axiom. **Bregman-Fenchel identification** (§6.3) is textbook Legendre-Fenchel; ties the segment’s divergence coordinate to #deriv-edge-update-natural-parameter’s update coordinate as Fenchel-dual aspects of one exponential-family structure. Fisher-metric-at-second-order is explicitly identified as *not* a distinguishing axiom within the f-divergence family (Eguchi’s theorem), closing the most-obvious alternative candidate uniqueness route as a no-go. **Information-theoretic-MDP lineage positioning** (§6.4) is *Discussion-grade* (positioning in the literature, not a derivation); its role is to place AAT’s direction-distinctive form honestly within the TP2011 / Rubin 2012 / Levine 2018 lineage rather than overclaim equivalence with any of them. The Rubin 2012 Theorem 3 PAC-Bayesian bound (§6.4) is a fourth independent operational motivation for the KL-to-reference form. The β_Σ naturalization (§7) is a trade-off analysis, not a derivation. The extensions (§9) are sketches, not derivations.

Max attainable: *exact* for (a) “forward-KL is a vacuous bound on R under deterministic π^* ” (direct calculation; §5); (b) “under deterministic π^* , $D_{\text{KL}}(\pi^* \| Q) = -\log(1 - \text{TV}(\pi^*, Q))$ ” (BH identity; §4); (c) “regret is Lipschitz-equivalent to $(1 - e^{-D_{\text{KL}}})$ with constants $\Delta_{\min} / V_{\max}$ ” (§4 matched lower bound); (d) “reverse-KL is the Bregman divergence of negative-entropy; log-odds is its Fenchel-dual natural coordinate” (standard Legendre-Fenchel; §6.3); (e) “reverse-KL is the unique chain-rule-respecting f-divergence in the direction-forced family” (Hobson 1969 / Csiszár 1991; §6.1); (f) “asymmetry is forced by regret’s one-sidedness” (direct; §5). *Robust qualitative* for the overall strengthened characterization — the chain-rule uniqueness is conditional on the chain-rule axiom, which is a principled axiom choice (AAT-internally motivated) but remains an axiom rather than a consequence of prior AAT commitments. The information-theoretic-MDP lineage positioning (§6.4) is *Discussion-grade* and honest about where AAT sits in the broader lineage rather than claiming formal equivalence.

Discussion.

The convergence with active inference’s reverse-KL. Active inference’s expected free energy (Friston et al. 2017; Parr & Pezzulo 2022) also uses reverse-KL against a preferred distribution, with the direction arising from the variational-free-energy-gradient lineage. AAT derives the direction from a regret-bound argument internal to the decision-theoretic scope. The convergence (two distinct derivations reach the same direction) is worth noting but is not a uniqueness argument — it shows reverse-KL is *natural* in multiple frames. The lineage difference preserves AAT’s distinctive content: AAT’s derivation does not require VFE as master objective or preferences-as-priors (7.5, 8.4 for the contrast).

Why the regret bound is the right derivation, not Donsker-Varadhan. The Donsker-Varadhan variational representation

$$\log \mathbb{E}_Q[e^f] \geq \mathbb{E}_P[f] - D_{\text{KL}}(P\|Q) \quad (\text{T.20})$$

gives reverse-KL bounds for exponential-moment functionals, but requires V to have bounded exponential moments (a stronger-than-bounded assumption for unbounded V). The Pinsker-TV route used here requires only bounded value range and is closer to the AAT scope’s natural assumptions. Donsker-Varadhan becomes relevant if unbounded- V extensions are attempted.

Relationship to standard policy-gradient regret bounds. The argument here is structurally identical to standard policy-gradient / online MDP regret bounds (Agarwal, Kakade, Lee, Mahajan 2021 “On the theory of policy gradient methods” Lemma 3.2; Even-Dar, Kakade, Mansour 2009 “Online Markov decision processes”). AAT’s contribution is *re-using the standard bound as the derivation of the KL direction in the strategy-cost objective*, not inventing the bound. This is consistent with AAT’s integration-over-invention character.

Position in the additive-coordinate-forcing pattern. The chain-rule uniqueness theorem of §6.1 is the *divergence-level* instance of an AAT-wide pattern in which logarithmic coordinates are uniquely forced via Cauchy’s functional equation on an AAT-internally-motivated additivity axiom. The pattern has two other primary instances: the chain-layer identity (8.1, the motivating anchor) and the update-layer log-odds uniqueness (U). 5.5 catalogs the pattern, states the 1-anchor-plus-2-theorem structure precisely, and documents Lyapunov and IB Lagrangian as adjacent family members that share the additive-decomposition shape without the AAT-internal forcing structure.

Working Notes.

- **Linear vs. square-root form — revisit with downstream evidence.** The segment keeps the linear-KL form for IB-shape alignment. If downstream work shows the IB-shape-across-four-operations claim in AA is strained and the four β s want separate scales anyway, switching 9.9 to the square-root form (naturalizing β_Σ) becomes attractive. Revisit with the O-BP2 (four compressions as one hierarchy) work if pursued.
- **Stochastic- π^* extension.** Flagged in §9 but not derived. The tied-optimum case is immediate; softmax-smoothed π^* needs a careful treatment of what regret quantity is being bounded (expected-vs-max).
- **Family characterization strengthening.** *Resolved 2026-04-22* — see §6.1. The chain-rule additivity axiom (Hobson 1969; Csiszár 1991 Theorem 3 corollary and Theorem 5; standard functional-equation derivation per Aczél & Daróczy 1975) picks reverse-KL uniquely within the direction-forced f-divergence family. The Fisher-metric candidate was ruled out by Eguchi’s theorem (every smooth f-divergence induces the Fisher metric at second order up to scale; not distinguishing). Full reasoning trail in `spikes/.integrated/spike-reverse-kl-uniqueness.md`, including the citation audit that corrected the initial draft’s misattributions (Csiszár 1972, Amari 2009 Theorem 1, and Amari & Cichocki 2010 Prop 3.2 do not contain the chain-rule uniqueness theorem; the correct attributions are above).
- **BH identity under deterministic π^* — elementary fact the earlier presentation underplayed.** *Resolved 2026-04-24.* §4 now presents the BH identity as primary, with Pinsker demoted to §4.1 as the loose general form. The earlier framing “non-Pinsker tighter bounds... BH flagged for operational work where the tightness difference matters” under-sold the result:

under AAT’s canonical deterministic- π^* scope, the BH inequality specializes to an *identity* (not merely a tighter inequality), making the KL-coordinate bound exact rather than approximate. The segment now carries the exact two-sided relation with matching lower bound via $\Delta_{\min}$.

- **Citation verification (Path 1 content, §6.4).** The information-theoretic-MDP lineage content in §6.4 was PDF-verified 2026-04-24 against TP2011 / Rubin-Shamir-Tishby 2012 / Levine 2018. The verification process found a specific attribution error in an earlier spike draft — the objective form $\max_{\pi} \mathbb{E}[R] - \beta^{-1} I(\Pi; X)$ with KL-to-reference reference is **Rubin 2012’s**, not TP2011’s (TP2011’s actual quantity is Information-to-Go multi-information $\mathfrak{I}^{\pi}$, a different object). The segment reflects the corrected attributions with explicit equation/page references. Full trail in `spikes/.integrated/spike-ib-purity-strategy-cost-strengthening-2026-04-24.md` §12 “Citation-verification addendum.”
 - **AAT-outlier-on-KL-direction — distinctive feature, not a departure.** *Documented 2026-04-24.* The literature uses agent-first or proposal-first directions; AAT uses optimum-first, forced by the regret-bound derivation and required for the BH identity to be exact. Owned rather than inherited; see §6.4.
 - **Cross-reference to NeurIPS Paper 2.** The deterministic- π^* corner identity is sharpened in NeurIPS 2026 Paper 2 (“A Unified Convergence Theory for Non-Stationary Reinforcement Learning”, `~/src/neurips/02-unified-convergence-rl/`, §5 Key Lemma 1 / `#lem-pointmass-identity`): under deterministic optimum, $\text{TV}(\delta_{a^*}, Q) = 1 - e^{-D_{\text{KL}}(\delta_{a^*} \| Q)}$ holds as an *exact algebraic identity*, lying strictly below the Pinsker and Bretagnolle-Huber envelopes at this corner (numerical comparison in Paper §App-E). The identity is **coordinate-optimal among bounds depending only on TV** — Paper 2 uses it as Component 2 of the composition theorem and as the structural reason why the segment’s optimum-first reverse-KL direction is the right coordinate. The classical BH bound is recovered as the loose general form away from the deterministic- π^* corner. See `spikes/neurips-back-integration-2026-05-08.md` §1 Paper 2 entry 1; standalone-segment proposal under back-integration §3 (proposed slug `#deriv-bh-pointmass-identity` or `#result-pointmass-tv-kl-identity`).
 - **Fenchel-Bregman connection to `#deriv-edge-update-natural-parameter`.** *Documented 2026-04-24.* §6.3 identifies reverse-KL and log-odds as Fenchel-dual coordinates of one exponential-family Legendre-Fenchel structure. The full geometric-unification reading (covering all four layers of `#disc-additive-coordinate-forcing`) is a meta-segment-level observation; the absorbing-spike trail is recorded in the spike-routing cycle (CHANGELOG 2026-05-17).
-

Appendix U

Derivation: Log-Odds as the Unique Additive-Evidence Parameterization for Edge Credences

The log-odds coordinate $\lambda_{ij} = \log(p_{ij}/(1 - p_{ij}))$ is the unique parameterization (up to positive affine transformation) on which independent Bernoulli evidence updates edge credences additively, under an evidential-additivity axiom motivated as the update-level analog of 8.1's chain-level additive log-confidence decomposition. This segment states the uniqueness theorem, derives it, and explains how it positions log-odds as the natural parameterization for AAT's continuous-gradient edge-update machinery.

This is the *update-level* instance of #disc-additive-coordinate-forcing's meta-pattern — the third formal layer where AAT commits to a specific coordinate because an internal additivity axiom forces it through a classical uniqueness theorem. Two prior layers compose with it: **log-probability at the chain layer** (chain rule of probability — #der-chain-confidence-decay); **reverse-KL at the divergence layer** (chain-rule additivity axiom via Hobson-Csiszár-Shore-Johnson — #deriv-strategy-cost-regret-bound); **log-odds at the update layer** (evidential-additivity axiom via Aczél's Cauchy functional equation — this segment). The Fisher information metric at the metric layer (Čencov's uniqueness theorem under parameterization-invariance) rounds out a four-layer structure that all live on the same geometric object — the exponential-family Legendre-Fenchel duality. *AAT's coordinates are not chosen by convention; they are forced by AAT-internal axioms via classical uniqueness theorems, and they all live on the same geometric object.*

The framework lifts this to a first-class coordinate commitment at the update layer. Continuous-gradient edge-update machinery in the credit-assignment treatment operates natively in log-odds and projects back to probability space via the sigmoid at the readout interface — the reason is mechanical: in probability space, the same update written as residual-times-gain can push the credence outside $[0, 1]$ when the Jacobian norm is small (a mechanical break of the probability domain). The log-odds coordinate has domain over all of $\mathbb{R}$, so additive updates never escape the domain; the sigmoid projection guarantees boundedness by construction, not by clipping. A downstream load-bearing consequence: the log-odds coordinate is what makes the detection-latency floor a structural rather than parametric result — the expected cycles to detect a regime change of given footprint scale linearly with accumulated experience precisely because the log-odds accumulation rate is forced (#deriv-update-detection-latency).

Formal Expression.

Setup. Let $p \in (0, 1)$ denote a scalar Bernoulli credence (the probability that a proposition is true) and let $\psi : (0, 1) \rightarrow \mathbb{R}$ be a smooth, strictly monotone reparameterization. Consider a sequence of independent Bernoulli observations $y_1, \dots, y_n \in \{0, 1\}$ drawn from a channel with likelihood ratio $P(y | H_1)/P(y | H_0)$.

Evidential-additivity axiom. The posterior update, applied to a single observation y , takes the form

Assumption (evidential-additivity axiom).

$$\psi(p_{\text{post}}) = \psi(p_{\text{prior}}) + g(y) \quad (\text{U.1})$$

for some function $g : \{0, 1\} \rightarrow \mathbb{R}$ that depends only on the observation y — not on p_{prior} nor on observation history.

Theorem.

Theorem. The functional equation above admits solutions if and only if

Derived (evidential-additivity uniqueness of log-odds, conditional on the axiom above).

$$\psi(p) = c \cdot \log \frac{p}{1-p} + d \quad (\text{U.2})$$

for constants $c > 0$ and $d \in \mathbb{R}$, with $g(y) = c \cdot \ell(y)$ where $\ell(y) = \log[P(y | H_1)/P(y | H_0)]$ is the log-likelihood ratio.

Derivation.

By Bayes' theorem applied to binary hypotheses,

Derived (Proof Step: Bayesian form of the update).

$$\frac{p_{\text{post}}}{1 - p_{\text{post}}} = \frac{p_{\text{prior}}}{1 - p_{\text{prior}}} \cdot \frac{P(y | H_1)}{P(y | H_0)} \quad (\text{U.3})$$

Taking the logarithm of both sides and writing $h(p) := \log(p/(1 - p))$,

$$h(p_{\text{post}}) = h(p_{\text{prior}}) + \ell(y) \quad (\text{U.4})$$

So $\psi = h$ trivially satisfies the axiom with $g = \ell$.

Suppose ψ is any smooth, strictly monotone reparameterization satisfying the axiom. Since the Bayesian mapping $p_{\text{prior}} \mapsto p_{\text{post}}$ is fully determined by y through the likelihood ratio, the difference $\psi(p_{\text{post}}) - \psi(p_{\text{prior}})$ depends only on y , and by the axiom must equal $g(y)$.

Derived (Proof Step: uniqueness by Cauchy functional equation).

Change variables via $\lambda = h(p) = \log(p/(1 - p))$ and define $\Psi(\lambda) := \psi(\sigma(\lambda))$ where $\sigma(\lambda) = 1/(1 + e^{-\lambda})$ is the logistic sigmoid. The axiom becomes

$$\Psi(\lambda + \ell(y)) - \Psi(\lambda) = g(y) \quad \text{for all } \lambda \in \mathbb{R}, y \in \{0, 1\} \quad (\text{U.5})$$

Extending to continuous-valued evidence (or considering mixtures of Bernoulli channels with varying likelihood ratios, which span all of $\mathbb{R}$ in the ℓ -value space), the identity

$$\Psi(\lambda + \ell) - \Psi(\lambda) = G(\ell) \quad \text{for all } \lambda, \ell \in \mathbb{R} \quad (\text{U.6})$$

holds for a function G independent of λ . This is the Cauchy functional equation (translation-additivity). Combined with the smoothness assumption on ψ , the unique solution class is $\Psi(\lambda) = c \cdot \lambda + d$ for constants c and d (Aczél 1966, *Lectures on Functional Equations and Their Applications*, §2.1).

Strict monotonicity of ψ forces $c \neq 0$. Taking ψ to have the same monotonicity sense as $p \mapsto p$ (credence increasing with ψ), we need $c > 0$. Thus $\psi(p) = c \cdot h(p) + d = c \cdot \log(p/(1-p)) + d$, and $g(y) = c \cdot \ell(y)$.

Derived (Proof Step: determining the constants).

This completes the proof. $\square$

Three-Layer Parallel.

The evidential-additivity axiom is the update-level instance of an additive-decomposition principle that AAT has already committed to at two prior layers:

Discussion (three-layer additive decomposition).

Table U.1

<i>Layer</i>	<i>Quantity decomposed</i>	<i>Decomposition form</i>	<i>Source</i>
Chain level	Confidence along a causal chain	$\log P(\text{chain}) = \sum_i \log P(E_i \mid E_{<i})$	8.1
Divergence level	Mismatch between optimal and strategy policies	$D_{\text{KL}}(\pi^* \parallel Q_{\Sigma_t})$ decomposes additively across DAG layers along the optimal trajectory	T §6.1
Update level	Credence evolution under independent evidence	$\psi(p_{\text{post}}) - \psi(p_{\text{prior}}) = g(y_1) + \cdots + g(y_n)$ for n observations, with $\psi = \log\text{-odds}$	This segment

Each layer forces a logarithmic coordinate through essentially the same structural move: products of independent factors become sums on a log scale. At the chain level, $p^n \rightarrow n \log p$; at the divergence level, $\prod Q \rightarrow \sum \log Q$; at the update level, $\prod \text{LR} \rightarrow \sum \log \text{LR}$. The three are the same transform applied to different quantities.

Interpretation for the Edge-Update Machinery.

For edge credence p_{ij} with log-odds $\lambda_{ij} = \log(p_{ij}/(1-p_{ij}))$, the Bayesian update under independent Bernoulli evidence is

Discussion (operational consequence).

$$\lambda_{ij}^{\text{post}} = \lambda_{ij}^{\text{prior}} + \ell(y) \quad (\text{U.7})$$

where $\ell(y)$ is the per-observation log-likelihood ratio.

Two operational consequences that follow from the uniqueness theorem:

1. **Domain unboundedness.** The log-odds coordinate has domain $\mathbb{R}$, not $[0, 1]$. Additive updates cannot escape the domain, regardless of update magnitude. The probability-space presentation $p_{ij} \in [0, 1]$ is the projected image of the log-odds coordinate, obtained via $p_{ij} = \sigma(\lambda_{ij})$ at the readout interface.
2. **Invariance under the chain of causal reasoning.** Because the log-odds coordinate is the unique additive evidence coordinate, evidence accumulated along one edge in a strategy DAG can be composed with evidence accumulated along another edge by addition in the log-odds vector space, provided the evidence is conditionally independent. The Beta-Bernoulli moment-parameter form $\hat{p} = \alpha/(\alpha + \beta)$ is the projected image, where α, β are the cumulative sufficient statistics in exponential-family form.

These consequences are why the continuous-gradient edge-update machinery in 9.6 is well-posed globally in log-odds but exhibits a mechanical break (unbounded updates pushing credences outside $[0, 1]$) when stated directly in probability space.

Scope Condition.

The evidential-additivity axiom applies to agent classes that treat observations as independent Bernoulli likelihood evidence — the Bayesian-coherent sub-scope of AAT. Non-Bayesian agents (PID controllers, rule-based systems, human judgment per 3.6) do not invoke likelihood-ratio accumulation and are outside the axiom's scope. This matches the sub-scope α / sub-scope β partition in 4.3: the uniqueness applies within sub-scope α , where B1 (directional fidelity) is already derived from Bayesian coherence.

Scope (evidential-additivity scope).

For multinomial / categorical edge credences with $K > 2$ outcomes, the analog is softmax / canonical exponential-family parameters: the softmax natural parameters $\eta_k = \log \pi_k$ (up to a reference-class shift) satisfy the same evidential-additivity axiom. The Bernoulli case ($K = 2$) collapses to log-odds.

Epistemic Status.

Derived (conditional on the evidential-additivity axiom). The uniqueness theorem is a standard functional-equation result (Cauchy functional equation with smoothness; Aczél 1966 §2.1). The AAT-internal motivation of the axiom is structural: it is the update-level analog of 8.1's chain-level additive log-confidence decomposition and T §6.1's divergence-level chain-rule additivity. Without the axiom, the selection of log-odds weakens to canonical-not-unique on convergent grounds (exponential-family naturalness, Fisher-information / natural-gradient canonicity, domain-well-posedness of the continuous gradient).

Max attainable: *exact* for the Cauchy functional equation step (standard); *derived-conditional* for the overall uniqueness claim because the axiom, though AAT-internally motivated, is itself a commitment rather than a consequence of prior AAT commitments. The axiom's status is parallel to the chain-rule additivity axiom in T §6.1 — both are AAT-internally motivated as the “right level” instances of a single additive-decomposition principle.

Not uniqueness in the unconditional sense. The theorem states that *if* credence updates live on a single additive coordinate whose increment depends only on the observation, *then* that coordinate is log-odds up to positive affine. It does not rule out non-additive update schemes, nor schemes that live on multiple coordinates (e.g., Beta-Bernoulli with sufficient statistics (α, β) rather than a single credence coordinate). The Beta-Bernoulli presentation is *equivalent* in content to the log-odds presentation under a change of coordinates; the uniqueness is about what single-coordinate additive form is possible.

Discussion.

Why log-odds is forced, not merely canonical. Three convergent grounds — exponential-family canonicity, Fisher / natural-gradient canonicity, and sector-condition preservation — establish that log-odds reparameterization is a natural coordinate for edge updates, but none of them is a uniqueness theorem. This segment supplies the uniqueness result conditional on the evidential-additivity axiom. The axiom is AAT-internally motivated, so the result is genuine strengthening — not “log-odds happens to work” but “log-odds is forced by the update-level analog of a principle AAT already relies on at chain and divergence levels.”

Why this matters for the mechanical break. The mechanical break in 9.6’s default signal function — unbounded gradient updates pushing credences outside $[0, 1]$ when $\|\mathbf{J}\|^2 \rightarrow 0$ — is a presentation artifact of the probability-space coordinate. In log-odds, the domain is $\mathbb{R}$ and the update is well-posed globally. The fix is not “clip the update” but “present the update in its native additive coordinate.”

Relationship to Bayesian conjugate analysis. The Beta-Bernoulli conjugate update produces posterior hyperparameters $(\alpha + y, \beta + 1 - y)$ that are additive in the sufficient statistics, giving the point estimate $\hat{p} = \alpha / (\alpha + \beta)$ that updates *non-additively* in probability space. The content is the same — the log-odds coordinate makes the additivity manifest at the level of credence itself, while Beta-Bernoulli shows the same additivity at the level of sufficient statistics. The two are dual presentations of the same exponential-family structure.

Non-uniqueness outside the Bayesian sub-scope. For non-Bayesian agents in sub-scope β of 4.3 (PID controllers, rule-based systems, human judgment), the evidential-additivity axiom does not apply. Such agents may update credences via coordinate-free heuristics (proportional blame, threshold rules) that do not decompose into likelihood-ratio addition. The uniqueness of log-odds is therefore scope-conditional, not universal.

Path B’s role in the overall pattern. The segment makes visible that AAT is committed to a family of consistent coordinate-forcing moves at chain, divergence, update, and metric layers, each forcing a specific coordinate via a uniqueness theorem on an AAT-internally-motivated axiom. Future coordinate-forcing structures in AAT (e.g., any novel decomposition along trajectories, across channels, or between agents) should be examined for whether Cauchy-FE on an additivity axiom or Čencov-invariance on a parameterization-invariance axiom applies. If yes, log-odds / natural-parameter / reverse-KL / Fisher-metric-style results compose cleanly. If no, the new layer is a genuine departure requiring its own analysis. The pattern is catalogued at the meta-pattern level in 5.5, including the 1-anchor-plus-3-theorem characterization (the chain layer is a mathematical identity; divergence, update, and metric are theorems conditional on AAT-internally-motivated axioms) and the adjacent-family classification (Lyapunov quadratic and IB Lagrangian share the additive shape but not the AAT-internal forcing structure).

Block-structured evidential additivity under L1’ correlated evidence. When sibling edges share a latent common cause C (L1’ correlated evidence per #def-strategy-dag’s Correlation Hierarchy), the independent-evidence axiom grounding Path B is literally false. Under *observable* C , the likelihood factorizes per-cluster as $P(\mathbf{y}, c) = P(c) \prod_k P(y_k | c)$, and the per-factor Aczél-1966 uniqueness argument applies independently to each factor. The forced coordinate becomes a $(2K + 1)$ -dimensional vector log-odds (one coordinate for θ_C , two per child edge for $p_{j|C}$ and $p_{j|\neg C}$), reducing to the scalar log-odds when the cluster is a singleton and no latent is present. This is a **generalization-in-scope** of the theorem above, not a new primary instance of #disc-additive-coordinate-forcing: the forcing machinery is unchanged (Cauchy-FE on per-factor additivity); the scope widens to block-factorized likelihoods under observable latents. Under *unobservable* C , the soft-EM responsibility-reweighted posterior depends nonlinearly on the prior, so no smooth ψ satisfies the block-additivity axiom — the axiom is structurally inconsistent with Bayesian mixture updates. This is the same scope boundary that #disc-identifiability-floor Instance 2 names via the Cramér-Rao Fisher-rank-1 obstruction; two independent analytical routes (Cauchy-FE failure at update layer; Cramér-Rao rank deficiency at observation layer) converge on the same unobservable- C structural floor, strengthening the “structural, not analytical-artifact” reading of Instance 2.

Findings.

Log-Odds as Uniquely-Forced Edge-Update Coordinate. Brief: When updating beliefs about whether something is true, there's a question of “what's the natural number to update?” — the probability itself, or some transformation of it. This result identifies a particular transformation (the log-odds, also used in logistic regression and Bayesian inference broadly) as the *uniquely* natural one for agents updating on independent evidence, in the precise sense that any equivalent representation must be log-odds up to scale. The choice isn't aesthetic; it's forced by what “independent evidence should add up” means mathematically. The uniqueness follows from Cauchy's functional equation operating on an evidential-additivity axiom motivated as the update-level analog of the chain-layer log-additive identity in `#der-chain-confidence-decay`.

Impact: Promotes a representational choice that previously had only convergent grounds (exponential-family canonicity, Fisher-information naturalness, sector-condition preservation) to a uniqueness theorem under an explicit AAT-internal axiom, joining the additive-coordinate-forcing meta-pattern (`#disc-additive-coordinate-forcing`) as one of its three theorem-grade instances. The result resolves a mechanical issue in the credit-assignment default signal function (`#disc-credit-assignment-boundary`), where unbounded gradient updates pushing credences outside $[0,1]$ in the probability-space presentation become globally well-posed in the log-odds presentation (domain $\mathbb{R}$). It also positions log-odds as the coordinate against which downstream gain-and-update machinery should be analyzed for sub-scope α (Bayesian-coherent) agents.

Novelty Claim: *Claim differentiation* on an already-canonical representational choice (log-odds as the natural Bayesian-update coordinate, well-known from logistic regression / exponential-family / information-geometry traditions) by deriving its uniqueness under an AAT-internally-motivated evidential-additivity axiom. The Cauchy-FE machinery is classical; the AAT-internal axiomatization and the meta-pattern membership are the contribution.

Related Work:

- Aczél 1966, *Lectures on Functional Equations and Their Applications* §2.1 — *formal antecedent* — Cauchy-FE uniqueness machinery directly adopted in the derivation; cited inline.
- Cox 1946, “Probability, frequency and reasonable expectation” *Am. J. Phys.* 14:1–13; Jaynes 2003, *Probability Theory: The Logic of Science* — *conceptual precursor* — log-odds and additive-evidence accumulation as the natural form for Bayesian update; folk knowledge in this tradition.
- Amari & Nagaoka 2000, *Methods of Information Geometry* §2 — *formal antecedent* — log-odds as the natural parameter in the Bernoulli exponential family; the canonical-natural-parameter framing.
- Bishop 2006, *Pattern Recognition and Machine Learning* §4.2; Murphy 2012, *Machine Learning: A Probabilistic Perspective* — *adjacent literature* — standard statistical-learning treatments of log-odds and logistic regression; representational standard.

Search Log: - 2026-04 (*intuition-only* on AAT-internal axiomatization angle): no targeted search for prior derivations of log-odds via the specific evidential-additivity axiom (update-level analog of chain-rule additivity) has been conducted. Pre-search expectation: the Cauchy-FE move is classical (Aczél 1966); the log-odds canonical parameter is exponential-family standard; the *AAT-internal axiomatization* (motivating the additivity axiom as the update-level analog of `#der-chain-confidence-decay`) is the contribution being claimed. A targeted search would query the Bayesian-coherence and information-geometry literatures specifically for prior axiomatic derivations of log-odds via Cauchy-FE in agent-update settings. - 2025 (*targeted*): Aczél 1966 confirmed as the formal antecedent for the Cauchy-FE machinery; cited inline.

Appendix V

Derivation: Adaptive-Gain Dynamics — A2’ Under a Learning Gain

AAT’s gain structure (3.6, 4.3) derives the optimal gain η^* per regime — the gain is a function of the noise model, chosen to minimize one-step mismatch variance. Real adaptive agents often *learn the noise model itself* (adaptive Kalman), *switch regimes* (IMM), *adapt step-size online* (RMSProp / Adam), or *optimize gain across tasks* (MAML). The gain becomes a state variable with its own update dynamics. The question is whether the sector-persistence machinery extends to this case, and where inside the A2’ sub-scope partition adaptive gain sits. The result: sub-scope α splits into α_1 (fixed-gain per Prop B.3) and α_2 (adaptive-gain under four derivable conditions named (MG-1)–(MG-4)), with sub-scope β catching the rest. The two-timescale argument is an augmented-state Lyapunov composition (standard Khalil Thm 4.18), not a Tikhonov reduction — the primary and meta-gain sector conditions compose rather than one being eliminated.

Four worked cases land the partition. **Adaptive Kalman** with a Mehra-type innovation-based estimator under window-length timescale separation derives at sub-scope α_2 from (MG-1)–(MG-4); under Mehra non-identifiability (rank-deficient transform matrix; Zagorbelny-Rawlins 2015, Dunik et al. 2021), (MG-2) fails structurally — a candidate instance of #disc-identifiability-floor on the meta-gain channel. **RMSProp** on strongly-convex losses is α_2 under design conditions (large β , coupling smallness); outside those conditions, vanilla Adam’s (MG-1) fails (Reddi et al. 2018 counterexample), and **AMSGrad’s monotonicity constraint on v_t is structurally a meta-gain repair that restores (MG-1) by construction**. **IMM regime transitions** sit at α_2 between transitions and β across the posterior-reconcentration window — explicit dwell-time machinery is required because (MG-1) fails uniformly in time during the transition. **MAML** has α_1 inner loops and a β outer loop, since the meta-loss is non-convex even under per-task convexity (Fallah et al. 2020) and (MG-2) is not derivable from per-task structure. Each placement is a derivation, not folklore: the choice between gain mechanisms has structural consequences for which sub-scope the agent’s persistence guarantee lives in, and (through the persistence-information-rate bound of #deriv-persistence-cost) for how much observation bandwidth the agent must consume to hold it.

Formal Expression.

Augmented-state setup. Treat the gain K_t as state. Define $\tilde{K}_t = K_t - K^*$ (error relative to a target optimal gain, specified per case: Riccati-steady-state for adaptive Kalman, EMA-fixed-point for RMSProp, etc.). The augmented state is $z_t = (\delta_t, \tilde{K}_t)$. Primary and meta-gain dynamics:

$$\dot{\delta} = -F(\delta; K^* + \tilde{K}) + w(t), \quad \dot{\tilde{K}} = -\Phi(\tilde{K}, \delta) + v(t) \quad (\text{V.1})$$

where F is the primary correction function (depending on the current gain via its argument), Φ is the gain-update contraction, w is primary-channel disturbance, v is gain-channel disturbance (estimator noise, innovation variability, etc.).

Meta-gain sector conditions (MG-1)–(MG-4).

(MG-1) Primary sector floor under bounded gain error. There exist $\underline{\alpha} > 0$ and $r_K > 0$ such that for all $\|\tilde{K}\| \leq r_K$ and $\|\delta\| \leq R$:

Formulation (meta-gain-conditions, extend A2' to adaptive-gain setting).

$$\delta^T F(\delta; K^* + \tilde{K}) \geq \underline{\alpha} \|\delta\|^2 \quad (\text{V.2})$$

— A2' uniform in the gain-error ball. The sector floor is preserved across the gain-state range the meta-learner visits.

(MG-2) Meta-gain sector condition. The gain-update map Φ satisfies (T1) (zero at $\tilde{K} = 0$) and a local sector bound:

$$\tilde{K}^T \Phi(\tilde{K}, \delta) \geq \alpha_K \|\tilde{K}\|^2 \quad \text{for } \|\tilde{K}\| \leq r_K, \text{ uniformly in } \|\delta\| \leq R \quad (\text{V.3})$$

with $\alpha_K > 0$. This is a sector condition in the gain-error state — the adaptive-gain analog of A2' itself.

(MG-3) Timescale separation. $\alpha_K \ll \underline{\alpha}$. The gain adapts slower than the primary state contracts. This is 4.9's convergence constraint transcribed onto Lyapunov decay rates instead of event rates.

(MG-4) Coupling boundedness. The gain-channel disturbance $v(t)$ has bounded contribution from the primary state:

$$\mathbb{E}[\|v(t)\|^2 \mid \delta] \leq \sigma_{K,0}^2 + c_v \|\delta\|^2 \quad (\text{V.4})$$

for some $c_v \geq 0$. (MG-4) with $c_v = 0$ is clean two-timescale decoupling; $c_v > 0$ is δ -coupled meta-gain disturbance (RMSPop near minimizer), requiring fixed-point closure.

Composed persistence result.

Under (MG-1)–(MG-4), the augmented state $z = (\delta, \tilde{K})$ is ultimately bounded in mean square. The Lyapunov candidate $V(z) = \frac{1}{2} \|\delta\|^2 + \frac{c}{2} \|\tilde{K}\|^2$ for appropriate weight c satisfies, along trajectories:

Derived (augmented-state-persistence, from sector-persistence-template applied twice with coupling).

$$\dot{V} \leq -\underline{\alpha} \|\delta\|^2 - c\alpha_K \|\tilde{K}\|^2 + \rho \|\delta\| + c(\sigma_{K,0} + \sqrt{c_v} \|\delta\|) \|\tilde{K}\| \quad (\text{V.5})$$

Complete-the-square on the cross term (requires $c\sqrt{c_v}$ small compared to $\underline{\alpha} \cdot c\alpha_K$, i.e., (MG-3) timescale separation plus (MG-4) coupling smallness):

$$\dot{V} \leq -\frac{\underline{\alpha}}{2} \|\delta\|^2 - \frac{c\alpha_K}{2} \|\tilde{K}\|^2 + \frac{\rho^2}{2\underline{\alpha}} + \frac{c\sigma_{K,0}^2}{2\alpha_K} \quad (\text{V.6})$$

Standard Lyapunov ultimate-boundedness (Khalil 2002 Thm 4.18) applied to the augmented state gives the composed persistence bound. Both δ and $\tilde{K}$ are ultimately bounded with explicit bounds in $(\underline{\alpha}, \alpha_K, \rho, \sigma_{K,0}, c_v)$. $\square$

A2' sub-scope partition: $\alpha_1 / \alpha_2 / \beta$.

The adaptive-gain analysis refines 4.2's A2' sub-scope partition into three tiers:

Formulation (sub-scope-refinement).

Sub-scope α_1 — fixed-gain, A2' derived. 4.3 Prop B.3's current scope: the gain K is treated as a static function of fixed noise model parameters. A2' is derived from B1 directional fidelity. Covers Kalman with known (Q, R) , conjugate-Bayesian updates, exponential-family MLE, linear correction with PD KH , strongly-convex-gradient fixed-step-size.

Sub-scope α_2 — adaptive-gain, A2' derived through augmented-state Lyapunov. When (MG-1)–(MG-4) hold with all four conditions derivable from the update-rule structure:

- Adaptive Kalman with Mehra-type innovation-based (Q, R) estimator under timescale separation
- RMSProp/Adam with strongly-convex loss, large β (slow EMA), and coupling-smallness; AMSGrad fix (Reddi et al. 2018) structurally a meta-gain repair preserving (MG-1)
- Any adaptive-gain scheme admitting a clean (MG-1)–(MG-4) derivation from the update rule

For these, A2' is derived at the augmented-state level, not merely assumed. Setting $\tilde{K} \equiv 0$ recovers sub-scope α_1 cleanly.

Sub-scope β — A2' assumed (possibly under scope narrowing). When any of (MG-1)–(MG-4) must be assumed per-agent rather than derived. Concrete instances:

- MAML outer loop: meta-loss non-convex even under per-task convexity (Fallah et al. 2020), so (MG-2) cannot be derived from per-task structure. Inner loop is α_1 ; outer loop is β .
- IMM regime transitions: (MG-1) fails uniformly in time during posterior-reconcentration windows of duration τ_{IMM} . Between-transition regime is α_2 ; across-transition window is β (scope narrowing via dwell-time + impulsive-disturbance absorption).
- Adam without AMSGrad on ill-conditioned problems: (MG-1) fails under aggressive- β + small-gradient-noise (Reddi et al. 2018 counterexample).
- Rule-based / PID / human-judgment adaptive gains: no structural argument for either (MG-1) or (MG-2); both must be assumed.

This refinement preserves the existing A2' sub-scope α as a specialization (α_1) and identifies a derivable adaptive-gain extension (α_2) with honest fallback to β when the derivability conditions fail.

Structured cases.

Case A — Adaptive Kalman with Mehra estimator. Scalar linear-Gaussian setting with unknown (Q^*, R^*) . The innovation-based Mehra estimator (Mehra 1970, 1972; Dunik et al. 2021 for identifiability) yields $\hat{Q}_t, \hat{R}_t$ from a sliding-window autocorrelation of the innovation sequence. For window length N , the gain-update map to first order in $\tilde{K}$ is a scalar Ornstein-Uhlenbeck process:

Derivation (case-adaptive-kalman-alpha2).

$$\tilde{K}_{t+1} = (1 - \lambda_N)\tilde{K}_t + \lambda_N\eta_t^{\text{inn}} \quad (\text{V.7})$$

with $\lambda_N \asymp 1/N$ (contraction rate set by window length) and η^{inn} zero-mean innovation noise with variance $O(1/N)$. This is itself an instance of **G** with:

- (T1): $\Phi_K(0) = 0$ — when gain is optimal, estimator returns optimal in expectation.

- (T2): $\tilde{K} \cdot \Phi_K(\tilde{K}) = \lambda_N \tilde{K}^2$, so $\alpha_K = \lambda_N$.
- (T3-S): bounded stochastic disturbance on the gain channel.

Prop A.1S applied to the meta-gain channel gives $R_{S,K}^* \asymp 1/\sqrt{N}$ — the classical Mehra asymptotic rate, now derived from (MG-2). Primary sector floor is preserved: $\underline{\alpha} = K^* - |\tilde{K}|_{\max}$ via triangle. Composed persistence via augmented-state Lyapunov gives $O(1/\sqrt{N})$ degradation from the fixed-gain case. **This case is derived at sub-scope α_2 .**

Under Mehra non-identifiability (rank-deficient transform matrix; see Zagrobelny-Rawlins 2015, Dunik et al. 2021), (MG-2) fails structurally — an instance of the 5.1 pattern on the meta-gain channel.

Case B — RMSProp on strongly-convex loss. The per-step effective step is $\eta_t^{\text{eff}} = \eta_t/(\sqrt{v_t} + \varepsilon)$ where $v_t = \beta v_{t-1} + (1 - \beta)\hat{g}_t^2$ tracks the second moment. Near the minimizer, $\mathbb{E}[\hat{g}_t^2] \rightarrow \|\nabla L\|^2 + \sigma_g^2$. Writing $\tilde{v}_t = v_t - \mathbb{E}[\hat{g}_t^2]$:

Derivation (case-rmsprop-alpha2-conditional).

$$\tilde{v}_{t+1} = \beta \tilde{v}_t + (1 - \beta)(\hat{g}_t^2 - \mathbb{E}[\hat{g}_t^2]) + \beta(\mathbb{E}[\hat{g}_{t-1}^2] - \mathbb{E}[\hat{g}_t^2]) \quad (\text{V.8})$$

The first two terms give (MG-2) with $\alpha_v = 1 - \beta$ and $\sigma_v^2 \asymp (1 - \beta)^2 \text{Var}(\hat{g}_t^2)$. The third is δ -coupled: $\mathbb{E}[\hat{g}_t^2]$ depends on δ through $\|\nabla L\|^2$, giving $c_v > 0$ in (MG-4).

Composed persistence under design conditions β close to 1 (slow EMA) and $\lambda_{\max}(H) \cdot R_S^* \ll \sqrt{\sigma_g^2}$ (coupling smallness): fixed-point closure between primary R_S^* and meta-gain R_v yields existence of a stable equilibrium. Sub-scope α_2 under design conditions.

Outside those conditions (aggressive β + ill-conditioning + small gradient noise — Reddi et al. 2018's Adam counterexample), fixed-point iteration diverges. **AMSGrad's monotonicity on v_t is structurally a meta-gain repair that restores (MG-1) by construction** — preserving sub-scope α_2 by forcing a condition (MG-1) the vanilla algorithm would violate.

Case C — IMM / regime-switching Kalman. Mixture of M Kalman filters with Markov-transition posterior over regimes. Between regime transitions, the posterior concentrates on the true regime and the effective gain approaches the regime-conditional optimum: sub-scope α_2 in the steady portion with Mehra-style derivation. Across regime transitions, posterior re-concentration takes $\tau_{\text{IMM}} \asymp 1/(1 - p)$ observations (self-loop probability p); during this window (MG-1) fails uniformly in time — the gain can be aligned with the wrong regime. Scope narrowing via dwell-time + impulsive-disturbance absorption: regime stable for $T_{\text{dwell}} \gg \tau_{\text{IMM}}$, transient mismatch bounded by counting-argument. Sub-scope α_2 between-transition; β across-transition.

Sketch (case-imm-alpha2-plus-dwelltime).

Case D — MAML inner-outer structure. Inner loop (per-task adaptation with k gradient steps) — sub-scope α_1 from Prop B.4 under per-task convexity. Outer loop (meta-parameter update via gradient on meta-loss $\sum_i L_i(\theta'_i(\theta))$) — Fallah et al. 2020's convergence analysis shows the meta-loss is non-convex even under per-task convexity because of the non-linearity of inner-loop updates in θ . (MG-2) is not derivable from per-task structure. Outer loop is β — A2' assumed per basin.

Classification (case-maml-mixed).

What Is Derived vs. What Is Chosen.

Table V.1

<i>Property</i>	<i>Source</i>	<i>Strength</i>
Augmented-state setup $z = (\delta, \tilde{K})$ with coupled sector dynamics	Definitional reformulation extending 3.2 + 3.6	Formulation choice
(MG-1) primary sector floor under bounded gain error	Derived in Cases A and B via triangle / ε -floor arguments	Derived (per case, conditional on regularity)
(MG-2) meta-gain sector condition	Derived in Cases A (Mehra OU) and B (EMA second-moment) from estimator structure	Derived (per named case)
(MG-3) timescale separation $\alpha_K \ll \underline{\alpha}$	Design condition on estimator window / EMA rate; Lyapunov decay-rate transcription of 4.9	Formulation (design condition; violations dissolve composition)
(MG-4) coupling boundedness of gain-channel disturbance	Decoupled ($c_v = 0$) in Case A; δ -coupled ($c_v > 0$) in Case B with fixed-point closure	Derived (per case)
Composed augmented-state persistence under (MG-1)–(MG-4)	Quadratic Lyapunov on z + Khalil 2002 Thm 4.18 ultimate-boundedness	Derived (conditional on MG-1 through MG-4)
Sub-scope refinement $\alpha_1 / \alpha_2 / \beta$	Extension of current A α/β partition	Formulation (classification)
α_2 reduces to α_1 under $\tilde{K} \equiv 0$	Direct substitution into augmented-state setup	Proved
Case A (adaptive Kalman): sub-scope α_2 under identifiability + window-length timescale separation	Mehra OU-form of estimator + Prop A.1S	Derived
Case B (RMSProp): sub-scope α_2 under design conditions; β otherwise; AMSGrad as meta-gain repair	EMA second-moment derivation + fixed-point closure + Reddi et al. 2018 counterexample	Derived (conditional); AMSGrad framing is discussion
Case C (IMM): sub-scope α_2 between-transition + β across-transition via dwell-time	Posterior re-concentration + impulsive-disturbance absorption	Sketch
Case D (MAML): inner-loop α_1 + outer-loop β	Fallah et al. 2020 meta-loss non-convexity	Classification (not derivation)
Mehra non-identifiability as meta-gain 5.1 instance	Rank-deficient transform matrix blocks (MG-2) derivation	Discussion (candidate floor instance)

Epistemic Status.

Conditional. Max attainable: *derived* for the composed persistence result under (MG-1)–(MG-4); *exact* for Case A (adaptive Kalman) under Mehra identifiability; *conditional* for Case B (RMSProp) under design conditions; *sketch* / *classification* for Cases C (IMM) and D (MAML).

The core claim — that (MG-1)–(MG-4) together give an augmented-state sector-persistence result via standard Khalil Thm 4.18 — is a clean application of two-timescale Lyapunov analysis to the gain-as-state formulation. The mathematical core is textbook (vector Lyapunov, augmented-state dissipation, completing-the-square on cross terms). The contribution is the AAT-framing: (i) naming the four conditions that make adaptive gain composable with sector-persistence, (ii) identifying which adaptive-gain schemes satisfy all four from update-rule structure (Cases A, B-design) and which require assumption (Cases C-transition, D-outer, rule-based), and (iii) refining the A2' sub-scope partition into $\alpha_1/\alpha_2/\beta$ as a principled extension rather than overhaul.

Imported machinery, acknowledged. The underlying techniques — two-timescale stochastic approximation (Borkar 1997, 2008), adaptive filtering (Mehra 1970–72; Mohamed-Schwarz 1999; Dunik et al. 2021), MRAC stability theory (Narendra-Annaswamy 1989; Ioannou-Sun 1996), non-convex meta-learning convergence (Fallah et al. 2020; Ji-Yang-Liang 2022) — are standard in control theory and ML. AAT's value-add here is the integration: connecting these techniques to G via the augmented-state Lyapunov composition, and producing a sub-scope partition that matches

AAT's scope-honesty architecture rather than importing the techniques' native terminology. The existing sub-scope α / β partition from A becomes α_1 / α_2 / β , with α_2 labeling the newly-derivable adaptive-gain region.

What this does not establish:

- Convergence rates beyond ultimate-boundedness (the bound gives stability, not speed).
- Behavior under (MG-3) violation — fast-adapting gain on a slow primary state — which is structurally different (the composite system becomes singular-perturbation pathological, not sector-persistent).
- Rigorous fixed-point closure for the Case-B δ -coupled case; the qualitative argument is clear but quantitative bounds on when the closure exists are per-problem.
- Composition at the multi-agent level — the augmented-state Lyapunov is for a single adaptive agent; team adaptive-gain (simultaneous meta-learning across agents) is beyond scope.

Discussion.

Resolving Epistemic Opacity. In #def-observation-function, the agent is structurally forbidden from knowing the true noise distribution ε_t . Yet the optimal gain #emp-update-gain requires knowing U_o . The adaptive Kalman filter with an innovation-based estimator (e.g., Mehra 1970) resolves this. The agent computes the autocorrelation of its observable mismatch sequence (innovations): $\iota_t = o_t - \hat{o}_t$. Because the innovation statistics are observable, the agent can estimate U_o (and U_M) purely from its own history $\mathcal{C}_t$, satisfying both the opacity axiom and the optimality requirement. The error in this estimation ($\bar{K}$) is exactly what the meta-gain sector condition (MG-2) bounds.

Why this is not a singular-perturbation reduction. Singular perturbation theory (Tikhonov 1952; Khalil 2002 Ch. 11) replaces fast dynamics with their quasi-steady-state manifold — the fast variable is *eliminated* in the slow reduction. The augmented-state analysis here does something different: it *keeps both states* and bounds their joint persistence via a weighted Lyapunov. Tikhonov gives “the fast variable tracks its slow manifold up to $O(\epsilon)$ ”; template-composition gives “each level's state is ultimately bounded with coupling-amplified disturbance” — the persistence-flavored state-ment aligned with AAT's other persistence results. They answer different questions; both are useful; neither replaces the other. M's stacking theorem (promoted 2026-06-10) is in fact stated in the template-composition register this paragraph describes — each level a template instance with coupling-amplified effective disturbance, both states kept, joint forward-invariance — with the quasi-steady-state manifold used only as the anchoring coordinate, not as an elimination step. The two segments now share the framing; what differs is the axis (gain-parameter timescale here, adaptive-level stacking there).

Three timescale patterns in AAT. M distinguishes model-class (slow) and parametric (fast) timescales. This segment introduces a *third* pattern: gain-parameter timescale, sitting between parametric-correction (fastest) and model-class (slowest). Adaptive gain is a different two-timescale pattern from structural adaptation — the gain is a parameter of the update rule, not of the model class. The convergence constraint $\nu_{n+1} \ll \nu_n$ from 4.9 translates in this segment to $\alpha_K \ll \underline{\alpha}$ on Lyapunov decay rates, which is the (MG-3) condition.

Meta-gain scope-honesty failure modes. - *Identifiability failures* on the meta-gain channel (Mehra non-identifiability) are instances of the 5.1 pattern: a structural no-go (non-identifiability) on the meta-gain channel blocks the sub-scope α_2 derivation. AMSGrad as a meta-gain repair is the constructive complement — an added structural condition (monotonicity on v_t) that restores (MG-1) and keeps the adaptive scheme in α_2 . - *Non-convex meta-losses* (MAML outer loop) put the outer loop in sub-scope β : A2' must be assumed within a basin, not derived from the meta-loss's global structure. This is the honest scope-narrowing posture — different from saying “it works” or “it doesn't” without qualification. - *Regime transitions* (IMM across switches) violate (MG-1) uniformly in time during posterior re-concentration windows. Repair via dwell-time is clean but significant scope narrowing.

Connection to 9.10's $\alpha_\Sigma = 1/(n+1)$. The strategy-edge persistence condition has α_Σ decaying with experience — this is adaptive gain of a specific form, where the gain (effective update rate per edge) is a function of accumulated pseudo-count. Without forgetting, $\alpha_\Sigma \rightarrow 0$ and (MG-1) fails asymptotically. 9.10's forgetting prerequisite $(1 - \lambda) > \rho_\Sigma/R_\Sigma$ is precisely the condition that keeps strategy-edge adaptation in sub-scope α_2 rather than letting it fall to β as n grows.

Relation to catastrophic forgetting via EWC. Elastic Weight Consolidation (Kirkpatrick et al. 2017) introduces a per-parameter stability-weighted update: gain is scaled by inverse Fisher information of prior tasks. In AAT vocabulary, EWC is a tensor-valued adaptive gain with stability weighting, distinct from the scalar adaptive-gain case treated here but fitting the same augmented-state framework with (MG-1)–(MG-4) reinterpreted per-parameter. Adapting the present derivation to EWC would tensorize the Lyapunov argument — a natural but not-yet-executed extension.

Working Notes.

- **Identifiability as meta-gain obstruction.** Mehra non-identifiability for certain system structures (Dunik et al. 2021) is an instance of 5.1 on the meta-gain channel — a structural no-go blocking (MG-2) derivation. Candidate for explicit cross-reference in 5.1's §"Adjacent Floors" as a fourth derived instance (after F1 CHT no-go, F13 Cramér-Rao mixture-identifiability, and Regime II-b misspecification-cost from 15.3).
 - **Rate-specific results.** The composed persistence result is a bound, not a rate. Can the $1/\sqrt{N}$ rate of Mehra-adaptive-Kalman be derived rigorously from the augmented-state Lyapunov? Classical Mehra asymptotics give it directly; deriving it from (MG-1)–(MG-4) would confirm the framework's completeness.
 - **Adversarial adaptive-gain.** If the environment adversarially varies K^* (hostile regime-switching cadence), dwell-time repair of Case C fails. Adversarial-tempo-advantage analog at the meta-gain level: faster meta-gain tracking beats faster environmental regime switching iff $\alpha_K T_{\text{dwell}}$ exceeds the transition rate. Not derived here; adjacent spike. (*Indexed: spikes/PROPOSED.md Tier 3 — "Adversarial adaptive-gain — meta-gain tempo-advantage analog"*)
 - **Discrete-time form.** Everything above is continuous-time. Discrete-time version involves the Lipschitz bound of $\mathbf{O}$ on both F and Φ . Extension is mechanical but deferred.
 - **Relationship to 9.10 time-varying α .** 9.10's $\alpha_\Sigma = 1/(n+1)$ decay is a concrete instance of adaptive-gain dynamics where (MG-1) fails asymptotically without forgetting. Worth re-reading that segment through the α_2 lens; the forgetting prerequisite is the specific move that keeps strategic adaptation in α_2 as n grows.
 - **Template-composition as general technique.** The augmented-state weighted-Lyapunov argument is a general two-timescale composition technique for **G**. It composes two template instances with a cross-coupling bound. A generalization to three or more timescales (inner loop + outer loop + meta-meta-loop, or fast + gain + slow + structural) would give a chain-composition form; not derived here but structurally available.
 - **AMSGrad as α_2 -preserving meta-gain repair.** Characterizing AMSGrad as "enforce (MG-1) by construction" is a clean AAT-framing of a pragmatically-motivated algorithm. Worth noting as an example of adaptive-gain schemes being AAT-classifiable by which of (MG-1)–(MG-4) they structurally ensure vs leave to luck.
 - **Deterministic-meta-gain special cases.** When the meta-gain is a deterministic function of the primary state rather than an independently learned variable — $K_t = \mathbf{I}^{-1}(\lambda_t)$ for Fisher whitening at the edge-update layer (**AM**), or $K = (H_M + H_L)^{-1}H_L$ for the Fisher-local model-parameter update (**AR**) — all four meta-gain conditions are satisfied trivially under the regularity conditions of those segments. The non-degenerate adaptive-Kalman-with-Mehra case is this segment's primary instance and adds a Mehra-style meta-channel with its own (MG-1)–(MG-4); the Fisher-local cases are degenerate special cases where the meta-gain is read off the primary state directly.
-

Appendix W

Derivation: Internal-External Decomposition of Agent Viability

Agent viability decomposes **additively** into agent-movable terms and an irreducible environmental floor — inherited from the exact mismatch decomposition — and provably does **not** factor multiplicatively into an agent-independent environmental rate times agent factors. The construction takes log-viability as the log-margin between steady-state mismatch and the critical task boundary, and inherits the additive split (estimation error + state-uncertainty floor + channel noise) from 3.5; the agent can move the estimation term by modeling and the state-uncertainty term by acting, the channel floor it cannot move at all.

The constitutive no-go is that multiplicative ρ -factorization of the form “ $\rho = \rho_{\text{external}} \cdot f(\text{model-class}) \cdot g(\text{policy})$ ” is type-incorrect: mismatch is constitutively agent-relative, so ρ^2 carries an estimation term that *vanishes* as the agent’s model improves — a contribution that subtracts to zero cannot be represented by attenuating a fixed agent-independent scalar by agent factors. Established exactly is the additive env/agent split with the environmental floor; the finer decomposition of the internal component into model, policy, and cross-term contributions is *conditional* and deliberately reserved — separating it requires Level-2 interventional access (5.1) and is under separate investigation.

Log-Viability.

Log-viability $\mathcal{V}$ is the margin by which steady-state mismatch R^* stays below the critical task boundary $\|\delta_{\text{critical}}\|$:

$$\mathcal{V} = \log \frac{\|\delta_{\text{critical}}\|}{R^*} = \log \|\delta_{\text{critical}}\| - \log R^* \quad (\text{W.1})$$

Persistence holds exactly when $\mathcal{V} > 0$. Under linear Model D, $R^* = \rho/\alpha$ (4.5), so

$$\mathcal{V} = \log \|\delta_{\text{critical}}\| - \log \rho + \log \alpha \quad (\text{W.2})$$

[*Derived*] — exact algebra (logarithm of a ratio); independent of how ρ itself is structured.

The disturbance is additive, not multiplicative.

The decomposition of viability into “what the environment imposes” versus “what the agent controls” runs through the structure of ρ . That structure is fixed by an existing exact result, not a free modeling choice.

The exact mismatch decomposition. By 3.5 (status: exact, via the fresh-noise assumption GA-1), expected squared mismatch splits with vanishing cross-terms:

Derived (exact, inherited from #result-mismatch-decomposition GA-1).

$$\mathbb{E}[\|\delta_t\|^2] = \underbrace{\mathbb{E}[\|\hat{o}_t - \hat{o}_t^B\|^2]}_{\text{estimation — reducible by modeling}} + \underbrace{\mathbb{E}[\text{Var}(\bar{o}_t \mid \mathcal{C}_{t-1}, a_{t-1})]}_{\text{state uncertainty — movable by acting}} + \underbrace{\mathbb{E}[\text{Var}(o_t \mid \Omega_t, a_{t-1})]}_{\text{channel noise — irreducible}} \quad (\text{W.3})$$

All cross-terms are zero by orthogonality (not independence), exactly as that result establishes. The first two terms are **agent-side**: the estimation term shrinks as the model improves, vanishing at the Bayes predictor; the state-uncertainty floor binds every model — well-specified or not — and yields only to *acting* (more informative interaction history), never to better modeling alone. The third is the **irreducible environmental floor**: a property of the observation channel the agent cannot reduce. A boundary note, because the working record has drawn it in two places: “the environmental floor” can be placed at the channel term alone (as here — everything else is agent-movable, by model or by policy) or at the full Bayes floor (state uncertainty + channel — everything no *model* can remove); the two placements differ by exactly the state-uncertainty term. This segment uses the first placement because the viability split is about what the agent — model and policy together — can move.

Rate-lift (stated convention). The persistence condition uses the disturbance *rate* ρ , related to expected squared mismatch by a fluid-limit identification $\rho^2 = \nu \cdot \mathbb{E}[\|\delta_t\|^2]$ for a positive throughput factor ν :

The viability decomposition. Substituting the exact split gives

Formulation — stated fluid-limit convention, introduced here; not a pre-existing canon identity. The structural conclusion below is convention-robust: any $\nu > 0$ preserves additivity, since positive scaling of a sum is a sum..

Derived (conditional on the stated rate-lift convention).

$$\mathcal{V} = \mathcal{V}_{\text{env}} + \mathcal{V}_{\text{agent}} \quad (\text{W.4})$$

where the first term $\mathcal{V}_{\text{env}}$ collects the channel-floor contribution (environmental affordance) and the second collects the estimation and state-uncertainty contributions together with the $\log \alpha$ tempo/gain terms (internal-operational health: the estimation part improves as the model improves; the state-uncertainty part as the agent’s interaction history grows more informative).

No multiplicative factorization (constitutive no-go).

There is no agent-independent scalar ρ_{external} with $\rho = \rho_{\text{external}} \cdot f(\mathcal{M}) \cdot g(\pi)$. Mismatch is constitutively agent-relative ($\delta_t \equiv o_t - \hat{o}_t$, 3.4), so ρ^2 carries the estimation term, which **vanishes as the agent’s model improves**. A multiplicative form attenuates a fixed agent-independent scalar by agent factors; it cannot represent a contribution that *subtracts to zero* under an agent-internal change. Multiplicative-in-rate and additive-in-(squared-mismatch) are type-incompatible, and $\sqrt{A+B} \neq \sqrt{A}\sqrt{B}$ blocks recovering the product form under the rate-lift. The decomposition is additive by construction, not by choice.

Derived — constitutive impossibility; exact under the constitutive definition of mismatch.

Epistemic Status.

The mismatch decomposition is **exact** (3.5, under GA-1). The viability decomposition in ρ -units is **conditional on the stated fluid-limit rate-lift** — a modeling convention introduced in this segment, not a pre-existing canonical identity; the *additive structure* (versus multiplicative) is robust to it for any positive throughput factor. The constitutive no-go is exact given the definition of mismatch. A general (non-isotropic, state/policy-dependent) diffusion $\Sigma(\delta, t)$ preserves additivity by linearity of the Itô generator in the diffusion matrix — this is an elementary extension *beyond* A Prop A.1S's constant-isotropic $\sigma_w^2 I_n$ case, correct but not itself part of that result's exact scope.

Discussion.

The finer split is conditional, not exact. Decomposing $\mathcal{V}_{\text{agent}}$ further into model-class ($\mathcal{M}$), policy (π), and cross contributions is a *mediation* question, not an identity: it is identifiable under interventional regimes (rotating the same agent across environments, or environments across agents) and **confounded** otherwise — in a feedback system high internal capacity itself lowers future environmental difficulty, entangling the terms. That confound is an identifiability-floor phenomenon (5.1): separating the agent-internal sub-terms from observational data alone is formally obstructed; Level-2 interventional access is required. This segment asserts only the exact two-term env/agent split and the no-go; the finer $\mathcal{M}/\pi$ /cross decomposition and its precise floor characterization are under separate investigation and are deliberately not asserted here.

Why this matters. The two-term split is the honest form of the internal/external diagnostic the persistence condition invites: it identifies exactly what an agent can improve (model error) versus the floor it cannot (channel noise), without the false promise that the environmental rate is an agent-independent quantity one can read off and attribute blame to.

Working Notes.

- Provenance (process history, not theory content): this segment previously asserted a multiplicative $\rho = \rho_{\text{external}} \cdot f(\mathcal{M}) \cdot g(\pi)$ split, was honesty-marked status: false, and was rebuilt on the exact mismatch decomposition. The constitutive no-go and the canon-forcing of the additive form were established by an independent recheck and confirmed by an independent verifier that also caught two pedigree over-claims (the ν rate-lift is *not* a pre-existing #hyp-mismatch-dynamics identity — landed here as a stated convention; the general- Σ additivity exceeds Prop A.1S's constant-isotropic scope — labeled as an elementary extension). Full trail: CHANGELOG 2026-05-18; spikes/ROUTING.md.
 - Open (reserved, under independent investigation): the conditional $\mathcal{M}/\pi$ /cross refinement and whether its Regime-C confound is a distinct #disc-identifiability-floor instance. Do not assert the three-way split until that resolves.
 - Boundary regression-guard (2026-07-03, audit 731548 B-2): the working layer twice defined the environmental floor $\rho_\star^2$ at *different* boundaries — the 2026-04-24 additive-variance spike at the full Bayes floor (state uncertainty + channel), the 2026-05-18 recheck at channel-instant variance only — differing by exactly the state-uncertainty term, and this segment's earlier gloss “vanishes for a well-specified model” was false under either placement (the Bayes-optimal predictor retains the state-uncertainty floor; see the Kalman anchor in 3.5). The three-term form now inherited from 3.5 states the boundary explicitly; if a future edit re-merges the middle term into either neighbor, check it against that segment's derivation first. History: CHANGELOG 2026-07-03; the ρ -family row in spikes/INDEX.md.
-

Appendix X

Derivation: Update Detection Latency Under Accumulated Experience

For a Beta-Bernoulli strategy-edge agent without forgetting, the expected number of cycles required to detect a within-class regime change of observable footprint ε scales as $\Omega((n_{\min} + 1)/\varepsilon)$ with $n_{\min}$ the minimum accumulated pseudo-count on load-bearing edges. The $1/(n + 1)$ rate is **structurally forced** — it is the log-odds update magnitude per cycle under $\mathbf{U}$'s Aczél-Cauchy-FE uniqueness theorem, and no choice of coordinate escapes it. The result sharpens 9.10's forgetting prerequisite from “required for asymptotic persistence” to “required for detection-latency bounded independently of operating point.” The broader myopia observation — that successful, high-capability organizations systematically underinvest in detecting regime changes that would require structural adaptation — admits a complementary decision-theoretic account via 10.2's oracle analysis; this segment's contribution is the signal-side lower bound, not the decision-side account.

Formal Expression.

Setup. An agent with a Beta-Bernoulli strategy DAG per 9.10 / $\mathbf{R}$, credit assignment via the log-odds signal of 9.6 (forced by $\mathbf{U}$), and no forgetting ($\lambda = 1$, pseudo-counts n_k accumulate monotonically). Let E_{load} denote the load-bearing edges on the current active plan; $n_{\min} = \min_{k \in E_{\text{load}}} n_k$. The agent has been operating with model-class fitness $\mathcal{F}(\mathcal{M})$ near 1, adaptive reserve $\Delta\rho^* > 0$, control regret δ_{regret} small — a high-operating-point configuration.

At cycle t_0 , a regime change occurs within the current model class (a true edge probability shifts by ε ; the L0 graph remains correct; the agent's model class still suffices). This is case R1 in the spike taxonomy — a within-class drift change. Other regime-change cases (R2 model-class inadequacy; R3 L0→L1 structural transition) are deferred to Working Notes and 5.1 respectively.

Detection-latency theorem.

Proposition. Under the setup, the expected number of cycles T_{detect} required for the log-odds coordinate on any load-bearing edge $k \in E_{\text{load}}$ to cross a fixed detection threshold $\Delta\lambda_{\text{detect}}$ in response to the regime change satisfies

Derived (detection-latency-R1, from edge-update-natural-parameter + deriv-edge-credence-dynamics).

$$\mathbb{E}[T_{\text{detect}}] = \Omega(\Delta\lambda_{\text{detect}} \cdot (n_{\min} + 1)/\varepsilon) \quad (\text{X.1})$$

Derivation. The default signal function of 9.6 (under the log-odds coordinate forced by $\mathbf{U}$) updates an edge's log-odds credence by

$$\lambda_k^{\text{new}} = \lambda_k + \eta_{\text{edge}} \cdot \iota_k \cdot J_k \cdot (y_G - \hat{R}_\Sigma) / \|\mathbf{J}\|^2 \quad (\text{X.2})$$

with $\eta_{\text{edge}} = 1/(n_k + 1)$ for Beta-Bernoulli (Prop B.4 of R). Per-cycle the expected log-odds update magnitude is bounded:

$$\mathbb{E}|\Delta\lambda_k| \leq \frac{|J_k| \cdot \mathbb{E}|y_G - \hat{R}_\Sigma|}{\|\mathbf{J}\|^2 \cdot (n_k + 1)} \quad (\text{X.3})$$

Under regime change R1 with observable footprint ε , the expected systematic residual $\mathbb{E}|y_G - \hat{R}_\Sigma| = \Theta(\varepsilon)$ as $\varepsilon \rightarrow 0$ (the misspecified-edge residual is proportional to the edge's probability shift, via the linearization of the Bernoulli likelihood in a neighborhood of the pre-change parameter). Combining: the per-cycle expected log-odds increment on load-bearing edges is $O(\varepsilon/(n_{\min} + 1))$.

For the agent to cross $\Delta\lambda_{\text{detect}}$, expected cycles required is at least $\Delta\lambda_{\text{detect}}/(O(\varepsilon/(n_{\min} + 1))) = \Omega(\Delta\lambda_{\text{detect}} \cdot (n_{\min} + 1)/\varepsilon)$. $\square$

The rate is structurally forced.

The $1/(n + 1)$ scaling is not a property of the specific update rule choice. Per U, the log-odds coordinate is the *unique* additive-evidence coordinate satisfying the evidential-additivity axiom — the Aczél 1966 Cauchy-functional-equation uniqueness theorem forces it up to positive affine transformation. In this forced coordinate, the per-cycle increment for Beta-Bernoulli edges is forced to be $O(1/(n + 1))$: the Fisher-equivalent statement holds in any sensible coordinate. Rearranging the update to a different scale does not change the rate — it just changes the units in which the rate is measured.

Observation (rate-forced-by-aczel).

The forcing composition:

1. U forces the log-odds coordinate via evidential additivity.
2. In that coordinate, Beta-Bernoulli accumulation gives $\eta_{\text{edge}} = 1/(n + 1)$.
3. Therefore the per-cycle update magnitude is structurally forced at $O(1/(n + 1))$.

This is the specific link between AAT's constructive meta-pattern (5.5, via U's theorem) and a downstream detection-latency consequence. The rate cannot be escaped without abandoning evidential additivity — which would invalidate the update rule on AAT-internal grounds, not merely operational ones.

Sharpening the forgetting prerequisite.

9.10 derives the forgetting prerequisite $(1 - \lambda) > \rho_\Sigma/R_\Sigma$ as required for *asymptotic persistence* — without forgetting, $\alpha_\Sigma = 1/(n + 1) \rightarrow 0$ and persistence eventually fails. The detection-latency theorem sharpens this to a **load-bearing claim about detection latency on the way to asymptotic failure**:

Corollary (forgetting-as-latency-bound, sharpens #schema-strategy-persistence).

Forgetting is required not only for asymptotic persistence, but also for detection latency to be bounded independently of operating point. Without forgetting, $n_{\min}$ grows monotonically, and $\mathbb{E}[T_{\text{detect}}]$ grows linearly with $n_{\min}$. With forgetting at rate $\lambda < 1$, the effective pseudo-count $n_{\text{eff}} = 1/(1 - \lambda)$ is bounded, and the detection latency caps at $\Omega(1/((1 - \lambda)\varepsilon))$ regardless of how long the agent has been operating.

This dualizes 9.10’s asymptotic claim: forgetting is operationally load-bearing at every step, not only in the limit. An agent with the right λ has bounded detection latency throughout its lifetime; an agent with $\lambda = 1$ has detection latency that grows unboundedly with experience, producing the phenomenon of stability-induced myopia in practice.

What Is Derived vs. What Is Chosen.

Table X.1

<i>Property</i>	<i>Source</i>	<i>Strength</i>
R1 detection-latency theorem $\mathbb{E}[T_{\text{detect}}] = \Omega((n_{\min} + 1)/\epsilon)$	U log-odds coordinate + R Prop B.4 $\eta_{\text{edge}} = 1/(n + 1)$ + Pinsker-type linearization	Derived (conditional on Beta-Bernoulli + log-odds + no forgetting)
$1/(n + 1)$ rate structurally forced	Composition of U’s Aczél-Cauchy-FE theorem with Beta-Bernoulli accumulation	Proved (conditional on evidential-additivity axiom)
Sharpening of forgetting prerequisite from asymptotic persistence to bounded detection latency	9.10 + this theorem	Derived
R2 model-class inadequacy sub-case (C1-diagnostic blindness under misspecification)	Common-mode-bias argument on $A_O^{(1)}$ vs $V_O(\pi_{\text{current}})$	Discussion (sketch; see Working Notes)
R3 L0→L1 structural transition sub-case	Direct application of 5.1 Instance 1 — already derived there	Reference (not new content)
Probability-of-deferral bound	Derivable from the expected-time bound via Markov inequality	Discussion
Connection to institutional inertia / Christensen / competency traps	Stability-induced-myopia as a mechanism for the empirical pattern	Discussion
Connection to IDT sidecar monitoring (Hafez et al. 2026 89% vs 44%)	On-policy reward-based monitoring is the log-odds accumulator; IDT bi-predictability is a change-detector that bypasses the $1/(n + 1)$ rate	Discussion

Epistemic Status.

Conditional. Max attainable: *exact* for R1 under stated scope (Beta-Bernoulli + log-odds forced + no forgetting + linearized residual near pre-change parameter); *robust qualitative* for general exponential-family edge with similar sufficient-statistic accumulation; *discussion-grade* for extensions beyond the Beta-Bernoulli / log-odds specialization.

The R1 theorem is a direct consequence of two derived AAT results (U’s forced log-odds coordinate + R’s Beta-Bernoulli η_{edge}) composed with a standard residual-magnitude argument (Pinsker-type inequality applied to the linearized Bernoulli likelihood). The mathematical core is not novel — what is novel is the AAT-framing: reading the $1/(n + 1)$ scaling as *structurally forced* through composition of existing AAT theorems, not contingent on update-rule choice. The rate is unescapable without abandoning evidential additivity.

The forgetting-prerequisite sharpening follows by direct substitution — replace unbounded $n_{\min}$ with bounded $n_{\text{eff}} = 1/(1 - \lambda)$ and the latency bound inherits the cap. This is where the segment’s most concrete contribution lies: the existing forgetting prerequisite’s justification expands from “asymptotic persistence fails without it” to “detection latency is unbounded in operating point without it.” Two load-bearing consequences instead of one, from the same $\lambda < 1$ condition.

What this does not establish:

- R2 (model-class inadequacy) as a full derivation. The sketch argument — that $A_O^{(1)}$ and $V_O(\pi_{\text{current}})$ share common-mode bias under misspecification, so the C1 regret diagnostic under-reports — is qualitatively clear but quantitatively informal. Making it rigorous requires a model of the bias’s functional form.
- R3 (L0→L1 transition) as new content. This is 5.1 Instance 1 — a structural no-go, not a latency bound. This segment references it but does not extend it.
- Decision-theoretic consequences. The theorem is about signal latency, not about whether the signal once received yields the right decision. 10.2’s oracle analysis (which shows deliberation rarely chosen) operates at the decision-theoretic level; the myopia theorem operates at the signal-quality level. An agent with degraded signals + perfect decision rule still defers; this is complementary to the oracle’s degraded-decision-with-perfect-signals result.
- Extension beyond Beta-Bernoulli. Gaussian edges admit $1/(n + c)$ scaling with analogous behavior; general exponential-family edges admit scaling following the sufficient-statistic accumulation. The qualitative form — expected detection time monotone in accumulated sufficient statistic — is more general than Beta-Bernoulli; the quantitative $1/(n + 1)$ is specific.

Discussion.

The forcing composition is where the meta-pattern interacts with downstream results. 5.5 is an AAT meta-segment naming the Cauchy-FE structural move at three layers (chain / divergence / update). This segment exhibits the specific case where that move’s downstream consequences bind: the log-odds forcing at the update layer implies — through Beta-Bernoulli accumulation — a specific detection rate that cannot be escaped via re-parameterization. The meta-pattern’s constructive power is visible here as a *negative* downstream consequence (bounded detection rate), which then becomes the positive motivator for forgetting. This is an example of AAT’s meta-architecture composing constructively: the forcing move (positive) produces a bound (negative) that sharpens an existing requirement (positive).

Distinction from adjacent phenomena:

- 4.8 characterizes WHEN restructuring is needed at truth (model-class inadequacy). This segment characterizes WHEN (in terms of cycles) the agent’s detection signal crosses threshold. The two together say: the agent will eventually detect a structural problem, but the latency can be large at high operating point.
- 10.2 gives a decision rule over exploit / explore / deliberate given signals. This segment is about signal quality — specifically, a predictable bias in signals as a function of operating point. The two compose: even an oracle’s decision rule applied to myopia-biased signals produces deferred action.
- 9.10’s forgetting prerequisite is a condition for asymptotic persistence. The corollary here is a condition for bounded detection latency — the same $\lambda < 1$ buys both.
- 5.1 Instance 1 (on-policy L0-insufficiency no-go) is a structural unboundedness result specific to L0/L1 transitions. The R1 detection-latency theorem is a quantitative bound on a broader class of regime changes (within-class drift, not just structural transitions).

Connection to cross-domain empirical patterns. Christensen’s *Innovator’s Dilemma*, Levitt & March competency traps, Hannan & Freeman structural inertia, March exploration/exploitation, and Eldredge-Gould punctuated equilibria all identify empirical patterns of stability-induced myopia across domains. The common story is that successful, mature, capable systems systematically underinvest in detecting regime changes that would require restructuring. The theorem supplies a mechanism: as experience accumulates (n_{min} grows), per-observation learning rate decays as $1/(n + 1)$, expected detection time grows linearly, and the structural-adaptation trigger signal is systematically delayed.

The phenomenon is not only decision-theoretic (incentive mis-alignment, short-horizon myopia) — it has a signal-side epistemic component forced by the log-odds coordinate. The signal-side and decision-side stories compose; neither subsumes the other.

Connection to IDT (Hafez et al. 2026) empirical result. Hafez et al. 2026 observed 89% vs 44% detection accuracy for Information Digital Twin (IDT) sidecar monitoring vs reward-based monitoring. The AAT-native explanation: reward-based monitoring is the log-odds accumulator — it inherits the $1/(n+1)$ detection rate and is subject to the theorem’s latency bound. IDT’s bi-predictability and entropy-change measures are structural-information-theoretic tests that bypass the log-odds coordinate entirely; they are change-detectors in the CUSUM / GLR tradition (Page 1954; generalized likelihood ratio tests) rather than Bayesian credence accumulators. The predicted gap between the two monitoring styles grows with accumulated experience — a testable refinement of the 89%/44% observation. A sidecar change-detector is always going to beat the log-odds accumulator on detection latency for within-class drift; the AAT theorem explains why without invoking the specific IDT mechanism.

The signal side and decision side as complementary. The fact that the oracle in 10.2 rarely chose deliberation is usually read as “agents correctly defer when deliberation is expensive and information is limited.” This segment’s theorem adds: “and the information is limited in a specific quantitative way that scales with operating point.” Both are true; neither is contradicted. The detection-latency result does not overturn the oracle analysis; it provides the signal-quality context in which that decision analysis operates.

Working Notes.

- **Migration note (2026-05-09 GUC rename):** Class 2 $\leftrightarrow$ Class 3 swap. Pre-2026-05-09: Class 2 = fully merged, Class 3 = partially modular. Post: Class 2 = Partial, Class 3 = Coupled. Removed at candidate stage per FORMAT.md Gate 4.
- **R2 sharpening (model-class inadequacy).** The argument: under C1 convention, $\delta_{\text{regret}}^{(1)} = A_O^{(1)} - V_O(\pi_{\text{current}})$ where both terms are computed under the current (now misspecified) M_t . Both terms bias in the same direction when the model class is structurally wrong; the gap stays small under common-mode bias. A clean version would compute the common-mode bias as a function of 3.5’s model-error component, show it is $O(1)$ (not $O(\epsilon)$) when the misspecification is directional, and derive a convention-dependent detection latency. Candidate one-evening spike. (*Indexed: spikes/PROPOSED.md Tier 3 — “R2 detection-latency sharpening under model-class inadequacy”*)
- **R3 reference, not extension.** The L0 $\rightarrow$ L1 structural-transition case is 5.1 Instance 1 — on-policy detection is forbidden by Bareinboim CHT, not merely slow. The detection-latency theorem here is about within-class drift (R1), which does not hit that no-go. When the regime change *does* involve an L0 $\rightarrow$ L1 transition, the theorem does not apply and the no-go does; the two coexist at different regime-change types.
- **Extension beyond Beta-Bernoulli.** Gaussian edges: $\eta \propto 1/(n+c)$ for the Gaussian-mean update. Exponential-family edges in natural-parameter form: η follows sufficient-statistic accumulation. The *qualitative* form (detection time monotone in accumulated sufficient statistic) is general; the *quantitative* $1/(n+1)$ is specific to Beta-Bernoulli. A parametric version of the theorem covering general exponential-family edges would be a clean follow-up.
- **Probability bound.** The expected-time bound implies a probability bound via Markov: $\Pr[\text{detect within } T] \leq T/\mathbb{E}[T_{\text{detect}}]$. Tighter concentration requires specifying the residual distribution’s shape; candidate one-paragraph extension.
- **Testable prediction.** In agents with controlled accumulated experience n , detection latency for regime change with fixed observable footprint ϵ should scale linearly in n . Directly testable in multi-arm bandits with drifting reward distributions, in simulated organizations with commit-log-mediated task environments, in human-subjects studies on noisy-signal regime detection. Each would confirm or refute the $1/(n+1)$ scaling in its native setting.
- **Multi-timescale connection.** The theorem adds empirical content to 4.9: the slow (structural-adaptation) timescale slows with accumulated fast-timescale state — the two timescales are not independent. The fast one’s accumulated state feeds back into the slow one’s responsiveness. This is the formal version of the institutional-inertia observation and is adjacent to, but not subsumed by, M’s current formulation.

- **Part III analog.** Do larger teams / organizations face proportionally longer detection latency via analogous accumulation? Candidate: team effective detection latency is bounded below by the slowest member's latency (the weakest-dimension version of 16.5). Connects to 15.1. Not derived here.
 - **Class 3 (Coupled) architectures.** For fully-coupled agents (LLMs), the Beta-Bernoulli accumulator is absorbed into the single merged computation — there is no literal edge-credence accumulator. But there *is* an analog at pretraining: token-statistic accumulation. Do LLMs exhibit a myopia analog? This is in the domain of 03-llm-core/, potentially connecting to #obs-context-turnover and the frozen-weights framing.
-

Appendix Y

Discussion: Independence Audit

AAT's results depend on a recurring modeling move: treat some quantity as independent of another to obtain tractable mathematics, then identify the failure regime where independence breaks and specify the repair. This segment enumerates the independence assumptions used across the theory, their failure regimes, their diagnostic signals, and the repair operations AAT provides. The enumeration makes visible what is *not* an independence assumption — acyclicity of Σ_t , Cox-derived probability, the Lyapunov machinery of G — and therefore what survives when a particular independence assumption fails.

Formal Expression.

Six load-bearing independence assumptions in AAT, each paired with its failure regime and repair:

Discussion (independence-audit).

1. Directed separation: M_t update independent of G_t . **Statement:** $f_M(M_{t-}, e_t)$ has no G_t argument — the epistemic update is goal-blind.

Where it appears: 6.3, the structural backbone of Part II. Feeds the orient cascade's sequential resolution (10.1), the causal validity of Q_O (6.6), and the scope of all Part II results to Class 1 (Separated) agents.

Failure regime: Class 3 (Coupled) architectures — transformer LLMs where attention processes goals and observations together. Motivated reasoning, confirmation bias, prompt-conditioned perception. Partially also Class 2 (Partial) agents.

Diagnostic signal: $\kappa_{\text{processing}} = I(G_t; M_{t+} \mid e_t, M_{t-}) / H(G_t \mid e_t, M_{t-})$. Zero for Separated agents; near one for Coupled; intermediate for Partial.

Repair operation: Class 2 (Partial) approximation quality scales with $\kappa_{\text{processing}}$. Class 3 (Coupled) agents require the coupled formulation $X_{t+} = f_X(X_{t-}, e_t)$ without (M_t, G_t) decomposition — the scope of 03-llm-core/. At the system level, Class 3 (Coupled) components can be wrapped in modular topology (separate observation processing, external monitoring — see #der-directed-separation Working Notes on the IDT pattern).

2. Causal sufficiency: no latent common causes among strategy nodes. **Statement:** Every common cause of two or more nodes in Σ_t is itself a node in Σ_t .

Where it appears: Q (precondition for the CMC-based Markov proof); 8.3 edge-independence in status propagation; all of Part II's strategy-layer formal results (Props B.1–B.6, the sector condition transfer of B.5, persistence of δ_s).

Failure regime: The dominant real-world case in complex, multi-stakeholder, or adversarial environments — shared infrastructure, market conditions, correlated adversary actions, common-mode risks, supply-chain dependencies.

Diagnostic signal: Pairwise covariance among sibling edges after edge credences have converged. Positive covariance rejects the independence hypothesis and localizes where a common cause is missing. See 9.2.

Repair operation: L1 augmentation — add common-cause nodes and restructure the DAG so each common cause is factored *above* the correlation it creates (8.3 Correlation Hierarchy; R Prop B.6). L0 formal results transfer exactly to correctly constructed L1 DAGs because L1 restores causal sufficiency by construction. The orient cascade's step 4c triggers this escalation.

3. Channel independence: observation channels contribute non-redundant correction. **Statement:** $\mathcal{T} = \sum_k \nu^{(k)}$. $\eta^{(k)*}$ with each channel contributing independently to effective tempo.

Where it appears: 3.8 (core definition); inherited by 4.5, 15.4, 15.1 (communication tempo), 13.2, and any result using scalar $\mathcal{T}$.

Failure regime: Any system with redundant or correlated information sources — overlapping sensors, correlated team-mate reports, redundant telemetry. In multi-agent settings, allies reporting the same intelligence source.

Diagnostic signal: $I(e^{(1)}; e^{(2)} \mid M_{t-})$ — pairwise mutual information between event streams conditioned on prior model state. Non-zero conditional mutual information *witnesses* dependence (a valid alarm), but it cannot serve as the correction term: it is sign-blind, and the sign of the noise correlation determines whether the additive form over- or under-counts (#deriv-tempo-additivity, no-go N1).

Repair operation: Under correlated channel noise the additive formula is *not* a general upper bound — the signed deviation $\Delta = \sum_k 1/U_o^{(k)} - \mathbf{1}^\top \Sigma_n^{-1} \mathbf{1}$ takes both signs (anti-correlated or heterogeneous positively-correlated noise makes the additive form an *undercount*; correlated noise is also a cancellation resource). The exact statements live in #deriv-tempo-additivity: additivity is exact under cross-channel noise independence (sufficient, not necessary), and in the common-source regime (shared persistent bias s) the additive form *is* an exact upper bound with closed-form penalty and joint information saturating at $1/\sigma_s^2$ — the echo-chamber bound. No channel-local additive correction exists in general for $n \geq 3$ (PID obstruction, no-go N2); precise tempo under correlation requires the joint Fisher form. Propagates into every segment using scalar tempo; per-segment ripple audit tracked in the 2026-07-15 spike.

4. Unity-dimension independence: ($U_M, U_O, U_\Sigma, U_{\text{obs}}$) substantially independent. **Statement:** The four unity dimensions of 14 parametrize composite quality along substantially independent axes.

Where it appears: 14 framing; 14.1's rate-distortion family indexed by each unity separately.

Failure regime: Shared models enable implicit strategic coordination (high $U_M \rightarrow$ high U_Σ without explicit policy sharing); aligned objectives often induce strategic coordination. Also *update-rule heterogeneity* (an axis not covered by any unity dimension) contributes to closure defect independently — the two-Kalman non-degenerate case shows $\varepsilon_x > 0$ from differing Kalman gains at perfect U_M .

Diagnostic signal: No clean one. The hypothesis is empirical and the theory currently logs it as working-position rather than resolved (see 14 Working Notes: update-heterogeneity gap).

Repair operation: Two options explored in the theory: (a) accept a two-axis structure (unity $\times$ homogeneity) for closure defect, (b) add a fifth update-homogeneity dimension. The current working position is (a); formal resolution open. Downstream uses of unity dimensions should not treat them as cleanly orthogonal.

5. Independent edge outcomes in AND/OR propagation. Statement: In $\hat{R}_\Sigma$ status propagation, sibling edge outcomes are independent given their parents.

Where it appears: 8.3 status propagation formula; 8.2; every downstream quantity computed from $\hat{R}_\Sigma$ (satisfaction gap, control regret via A_O , strategy-plan-confidence error δ_s).

Failure regime: Same as causal sufficiency (item 2) — the CMC theorem makes them the same condition: causal sufficiency $\Leftrightarrow$ exogenous noise independence $\Leftrightarrow$ edge-outcome independence. When one fails, they all fail.

Diagnostic signal: Same as item 2.

Repair operation: Same as item 2 (L1 augmentation). The assumption is not a separate modeling choice from causal sufficiency; the theory treats it as the *consequence* of causal sufficiency (via CMC), not an independent axiom.

6. Scalar tempo / isotropic correction. Statement: $\mathcal{T}$ as a scalar captures the agent's correction capacity.

Where it appears: 4.5 linear operational form; most Part I results stated in scalar $(\mathcal{T}, \rho, \|\delta_{\text{critical}}\|)$ form; 9.8 aggregate.

Failure regime: Any real multi-dimensional system with non-uniform correction capacity across dimensions. Simulation confirms scalar $\rho/\mathcal{T}$ overestimates by up to 72% in anisotropic systems, with the weak dimension accounting for 84% of total mismatch.

Diagnostic signal: Gain variation across dimensions (easily computable when gains are explicit — Kalman, gradient descent with diagonal preconditioning).

Repair operation: Per-dimension persistence: $\mathcal{T}_k > \rho_k / \|\delta_{\text{critical},k}\|$ for each dimension. See 16.5. The weakest dimension is the bottleneck. Same structural pattern propagates to strategic tempo (9.8 Per-edge persistence) — the bottleneck edge determines persistence, not the aggregate.

What is not an independence assumption.

Several results that might appear to depend on independence actually do not:

- **Acyclicity of Σ_t** is proved from temporal ordering over a finite horizon (Q). Independent of any independence assumption.
- **Cox-derived probability structure** at edges is independent of edge-independence — Cox forces the *measure* to be probability, not the joint structure.
- **Sector-persistence template machinery** (G) depends only on the sector condition (T2) and bounded disturbance (T3); it does not require any independence assumption on the state variable or correction function.
- **The DAG structure derivation** (directed edges, acyclicity, representability as a Bayesian network) does not require causal sufficiency — only the *Markov factorization* within the derived structure does.

This is the boundary between AAT's robust results and its conditional ones: robust results survive when the independence assumptions fail; conditional results break and require their specific repair operations.

Epistemic Status.

Robust-qualitative. Max attainable: *robust-qualitative*. The enumeration is definitional — each independence assumption is stated in a named segment, and the repair operations are either formally specified (causal sufficiency $\leftrightarrow$ L1 augmentation) or named open problems (unity-dimension independence). What this segment adds is the cross-cutting view: the six assumptions are the recurring modeling move that distinguishes AAT’s exact-tractable results from its conditional results, and the repairs are (mostly) already present in the theory but scattered across segments.

The enumeration is not exhaustive. Other independence-flavored assumptions exist in specific instantiations (e.g., independent Beta priors across edges in the credit-assignment analysis of 9.6). The six listed are the load-bearing ones — those whose failure regime the theory has actually characterized and provided a repair for.

Discussion.

Why enumerate these in one place. AAT’s theoretical contribution is often described as “integration”: bringing control theory, causal inference, information theory, and agent architecture under one framework. Integration *is* achieved largely through the independence assumptions listed above: each is a bridge that lets a module from one discipline plug into a module from another. Directed separation lets control-theoretic Lyapunov analysis operate on the epistemic substate without goal-entanglement. Causal sufficiency lets Pearl’s DAG machinery apply to strategy without requiring the agent to model every environmental common cause. Channel independence lets the tempo framework sum over heterogeneous observation modalities.

When these bridges fail, the integration does not fail catastrophically — each failure has a characterized repair. But the theory’s “exact” claims depend on the bridges holding, and recognizing this is the honest way to present the framework’s scope. A reader who understands the six assumptions and their repairs understands where AAT’s results apply exactly, where they apply approximately, and where they require structural revision.

Connection to Z. The L0/L1 distinction for causal sufficiency (items 2 and 5) is one instance of a meta-pattern: AAT’s results are parameterized by approximation level, with proved monotonicity between levels. See Z for the full treatment — the pattern also applies to the C1/C2/C3 value-convention hierarchy and the Tier 1/2/3 bridge-lemma contraction structure.

What this does NOT say. The enumeration does not claim the theory’s results are fragile or routinely fail. Most independence assumptions hold approximately in most well-behaved domains, and the theory’s quantitative predictions are typically within the regime where assumptions apply. The audit is a scope clarification, not a criticism.

Working Notes.

- **Migration note (2026-05-09 GUC rename):** Class 2 $\leftrightarrow$ Class 3 swap. Pre-2026-05-09: Class 2 = fully merged, Class 3 = partially modular. Post: Class 2 = Partial, Class 3 = Coupled. Failure-regime references under Assumption 1 updated accordingly. The C1/C2/C3 mention at line 120 is the value-convention hierarchy (a different ladder), not the GUC axis — left unchanged. Removed at candidate stage per FORMAT.md Gate 4.
- **Interaction between assumptions.** The six assumptions are not fully orthogonal. Causal sufficiency failure and edge-independence failure are the same failure (item 5 is a consequence of item 2 via CMC). Directed separation failure can propagate through the orient cascade into the strategy-layer results (which inherit their structure from the cascade). A fuller treatment would map the dependency graph among assumptions and characterize cascading failure modes.
- **Quantitative failure modes.** For each independence assumption, the theory has or could have a quantitative failure bound — how much the dependent result degrades as a function of the independence-violation magnitude ($\kappa_{\text{processing}}$, the covariance magnitude, the mutual information between channels). These bounds are partially specified in segment-level caveats; collecting them in one place would strengthen the audit.

- **Empirical audit.** A concrete agent could be scored against the six assumptions to produce an “independence profile.” For software agents: directed separation $\sim$ holds (Separated), causal sufficiency $\sim$ holds (test-based Regime A), channel independence $\sim$ holds (tests target disjoint behaviors), scalar tempo may fail (component-level gain variation). For LLM agents: directed separation fails structurally (Class 3 Coupled), causal sufficiency varies by task, channel independence varies. The profile would be a diagnostic for where the theory’s exact claims apply to the agent and where they degrade.
-

Appendix Z

Discussion: Approximation Tiering

AAT uses a recurring meta-pattern for handling intractability: when a problem admits no tractable exact treatment in general, introduce a tiered hierarchy of approximations with proved monotonicity between tiers and a diagnostic for ascending when needed. Three such hierarchies exist in the theory — the Correlation Hierarchy (L0/L1/L2) in 8.3, the Convention Hierarchy (C1/C2/C3) in 6.6, and the Tier 1/2/3 contraction taxonomy in 13.1. This segment articulates the pattern explicitly, identifies what makes a successful approximation tiering, and notes where other scattered simplifications might fit the same shape.

Formal Expression.

A successful approximation tiering has four components:

Discussion (approximation-tiering).

(AT1) **A parameter indexing tractability.** Some quantity (correlation modeling depth; continuation-policy fidelity; operator regularity) that admits a natural ordering from least to most demanding computation.

(AT2) **Monotonicity.** A proved ordering of the results produced at each level — each higher tier dominates the lower in the direction of interest (calibration accuracy, diagnostic force, guarantee strength).

(AT3) **Graceful degradation.** Lower tiers produce usable (not vacuous) results, with the specific form of the degradation characterized. An agent stuck at the lowest tier knows what it is missing, not just that it is incomplete.

(AT4) **Ascension diagnostic.** A signal the agent can detect at a lower tier that indicates escalation to a higher tier is warranted. Without this, the tiering is descriptive but not operational.

The three AAT tierings.

Table Z.1

<i>Hierarchy</i>	<i>Parameter</i>	<i>Monotonicity</i>	<i>Lowest tier</i>	<i>Highest tier</i>	<i>Ascension diagnostic</i>
Correlation (8.3)	Modeling depth of inter-edge dependence	L0 is conservative; L1 is unbiased on augmented graph; L2 is the full joint	L0 (independence model, $O(V + E)$ propagation)	L2 (full joint, $O(2^m)$ in general)	Sibling-edge covariance after credence convergence (9.2)

<i>Hierarchy</i>	<i>Parameter</i>	<i>Monotonicity</i>	<i>Lowest tier</i>	<i>Highest tier</i>	<i>Ascension diagnostic</i>
Convention (6.6)	Continuation-policy fidelity	$A_O^{(1)} \leq A_O^{RH} \leq A_O^B$ (proved monotonicity)	C1 (one-step improvement)	C3 (Bellman optimal)	Persistent $\delta_{\text{sat}} > 0$ under C1 that C2 replanning may resolve
Contraction (13.1)	Operator regularity for bridge lemma	Tier 1 $\supset$ Tier 2 $\supset$ Tier 3 in terms of proved contraction strength	Tier 3 (domain-specific verification required)	Tier 1 (contraction proved for full class)	Structural test: is the correction strongly monotone, locally convex, or neither?

Each row satisfies (AT1)–(AT4) fully — these are mature tierings in the theory. Each tiering is independently motivated by a different source of intractability, and each has its own ascension mechanism. Their structural similarity is not designed-in; it is the consequence of facing three separate intractable problems with the same modeling strategy.

Candidate future tierings. Several other scattered simplifications in the theory have the right shape to become tierings but are not currently organized that way:

- **Scalar vs. per-dimension tempo** (around 3.8, 16.5). Tier 0 (scalar isotropic), Tier 1 (per-dimension), Tier 2 (full tensor with off-diagonal coupling). Monotonicity: scalar overestimates margin, per-dimension is exact for diagonal-gain systems, full tensor handles off-diagonal gain coupling. Ascension diagnostic: gain variance across dimensions.
- **AND/OR parameterization** (8.2). Tier 0 (AND/OR, $O(k)$ parameters per node), Tier 1 (weighted/threshold, $O(k)$ additional parameters), Tier 2 (full CPT, $O(2^k)$ parameters). Monotonicity: expressiveness. Ascension diagnostic: systematic miscalibration at AND/OR nodes that survives L1 augmentation.
- **Identifiability regimes** (9.5). Regime A/B/C is already a tiering of causal identification quality, with t_{ij} as the indexing parameter. Monotonicity: interventional $\geq$ partial $\geq$ observational. Ascension diagnostic: available intervention opportunities in the domain.

Promoting these to named tierings with explicit (AT1)–(AT4) structure would make AAT’s scope parameterization more uniform. This is not pursued here — it is noted as a direction for editorial consolidation.

Epistemic Status.

Robust-qualitative. Max attainable: *robust-qualitative*. The meta-pattern is an editorial observation: three existing hierarchies in the theory share the same four-component structure, and the shared structure is worth naming because it makes the theory’s scope-parameterization move visible rather than implicit. The pattern itself is not a theorem — there is no result that says “every intractable AAT problem admits a (AT1)–(AT4) tiering.” The candidate future tierings are conjectures about where the pattern could be extended; each would require its own monotonicity proof and ascension diagnostic.

The three existing tierings are individually grounded: the L0→L1 monotonicity holds under the CMC theorem and augmentation construction; the C1→C2→C3 monotonicity is proved in 6.6; the Tier 1/2/3 taxonomy is grounded in operator theory (strong monotonicity, local convexity, arbitrary).

Discussion.

Why this matters. AAT's results are often presented as "exact under assumption X." A reader can misread this as "the theory works when X holds and breaks otherwise." The tiering pattern reveals a different structure: when X fails, a named lower tier takes over with a characterized weaker result, and the theory provides a diagnostic for when to escalate to reinstating X by structural repair. This is graceful degradation — a property that a theory wanting to apply across a range of agents and domains needs, and that AAT achieves in three places independently.

Reading AAT through the tiering lens: the theory's "exact" core is what it guarantees at the highest tier of each hierarchy. Its "conditional" periphery is where specific tiers are in force. The claim is not that agents should always operate at the highest tier — that would often be intractable — but that the theory tells the agent which tier it is in and how to get to a higher one when the binding constraint changes.

Connection to Y. The two meta-segments are complementary. Y enumerates the independence assumptions whose failure drops a result from exact to conditional. This segment enumerates the tierings that provide the structured *recovery* from those drops. Together they characterize AAT's scope parameterization: the independence audit says where the boundaries are; the tiering pattern says how to navigate within them.

What a full tiering promotion would look like. If the scattered simplifications listed above were promoted to named tierings, the theory's scope-parameterization would be dramatically more uniform: every conditional result would come with its own tiering and ascension diagnostic. This is aspirational. The three existing tierings took substantial segment-work to formalize (a full Appendix derivation for the CMC, a monotonicity proof for the convention hierarchy, a spike for the incremental sector bound). Extending the pattern to five or six hierarchies would be correspondingly substantial.

The pattern is not unique to AAT. Approximation tiering is a common move in applied mathematics — numerical methods (low-order vs. high-order integrators with error bounds), information theory (typical vs. exact coding with rate-distortion monotonicity), statistical inference (maximum-likelihood vs. Bayesian vs. nonparametric). AAT's contribution is not the tiering pattern itself but its deployment across the specific intractable problems that adaptive-agent theory faces.

Working Notes.

- **Formal characterization of the tiering pattern.** Is there a mathematical structure that unifies approximation tiering across domains? Candidates: lattices of model classes ordered by inclusion, rate-distortion curves with explicit corners, hierarchies of Galerkin approximations. None of these is an exact match; the AAT tierings are closer to "ordered families of sufficient conditions" than to any standard mathematical hierarchy. Worth investigating whether a cleaner abstract pattern exists.
 - **Interaction between tierings.** An agent operating at L1 correlation, C2 convention, and Tier 2 contraction is in a specific combined regime. The cross-hierarchy interactions are not worked out — is there cross-hierarchy monotonicity? (E.g., does L1 change anything about the convention hierarchy's guarantees?) Each tiering is independently-grounded but the joint structure is not yet mapped.
 - **Diagnostic costs.** Each ascension diagnostic has a cost: detecting sibling covariance requires observing convergence; detecting C1-inadequacy requires comparing replanning values; detecting Tier 1 violations requires verifying strong monotonicity. A unified treatment of "when is the diagnostic itself worth running?" would connect this segment to 4.1 and the allocation analysis in 10.2.
-

Appendix AA

Discussion: Compression Operations in AAT

AAT contains four compression operations — the epistemic model M_t , the strategy DAG Σ_t , shared intent G_t^{shared} , and the composition projection Λ — each formulated in its own segment with its own objective. Three of the four are written in Information Bottleneck (IB) form already; the fourth is stated as an IB constraint. This segment makes the shared shape explicit, promotes one underspecified source (the ontologically ambiguous “true causal structure” for Σ_t) to a cleaner formulation parallel to M_t , and establishes that composition admissibility (P1) is the Lagrangian-dual of a standard IB objective. It does *not* claim the four operations reduce to a single optimization problem — cross-instance theorems do not follow from the shared shape alone, and several conditions (Lipschitz regularity (P2), dimensional reduction (P3) in the Gaussian case, interventional relevance for Level-2 edges) remain outside the IB frame.

Formal Expression.

The shared IB shape.

Every compression operation in AAT has an objective or constraint of the form:

Discussion (ib-shape).

$$T^* = \arg \min_{T|X} [I(X; T) - \beta \cdot I(T; Y)] \quad (\text{AA.1})$$

with the Markov chain $Y - X - T$. The four AAT instances specialize this with different bindings:

Table AA.1

<i>Instance</i>	<i>X (source)</i>	<i>T (compressed)</i>	<i>Y (relevance variable)</i>	<i>β (trade-off)</i>
Model compression (2.2)	$\mathcal{C}_t$	M_t	$O_{t+1:\infty} \mid a_{t:\infty}$	$\beta(\rho, \pi)$ — volatility and policy
Strategy compression (9.9)	$\mathcal{C}_t$	Σ_t	$\pi^* \mid M_t$	β_Σ — cognitive cost per decision-bit
Shared intent (14.2)	$G_t^{\text{full}} = (O_t, \Sigma_t)$	G_t^{shared}	$a_t^{\text{coordinated}}$	bandwidth per coordination-bit

Instance	X (source)	T (compressed)	Y (relevance variable)	β (trade-off)
Composition projection (13.1 P1)	$X_{\text{micro},t}$	$\Lambda_X(X_{\text{micro},t})$	$o_{\text{micro},t+1} \mid a_{\text{micro},t}$	$\beta(\epsilon_I)$ — rate-distortion Lagrange multiplier

What the four instances share: *shape* (the objective structure), *variational calculus* (minimization over stochastic compressors), and *rate-distortion interpretation* (each trade-off parameter indexes a point on the frontier). What they do not share: source type (history vs. structured state), relevance-variable availability (observed vs. latent), computability (Gaussian closed forms vs. variational approximation), or a single joint optimization across the four. The level of unification is *medium*: shared shape and vocabulary, not a shared master problem.

Strategy compression: source reformulation.

The current statement in 9.9 has the Σ_t IB objective as:

Formulation (strategy-compression-source).

$$\Sigma_t^* = \arg \min_{\Sigma_t} [\text{DL}(\Sigma_t) - \beta_{\Sigma} \cdot I(\Sigma_t; \pi^* \mid M_t)] \quad (\text{AA.2})$$

Two issues with this as currently written:

1. **The compression cost is description length, not mutual information.** DL and $I(X; T)$ are related through coding-theoretic equivalences but coincide only under specific coding schemes. Using DL in the complexity term blocks the identification of this objective as an IB instance directly.
2. **The “source” is not an AAT object.** To fit the IB shape, the objective implicitly treats Σ_t as a compression of “the true causal structure.” That structure is not part of AAT’s ontology — the agent never has access to it; it is only ever implicit.

Reformulation. Treat Σ_t as a compression of $\mathcal{C}_t$ (the interaction history — the agent’s only evidence) *for decision-relevance*, parallel to M_t which is a compression of $\mathcal{C}_t$ *for prediction-relevance*:

$$\Sigma_t^* = \arg \min_{\Sigma_t} [I(\mathcal{C}_t; \Sigma_t) - \beta_{\Sigma} \cdot I(\Sigma_t; \pi^* \mid M_t)] \quad (\text{AA.3})$$

Under this reformulation:

- The source $\mathcal{C}_t$ is a well-defined AAT object, shared with the M_t instance.
- The two instances differ cleanly in relevance variable: M_t is compressed for prediction ($Y = o_{t+1:\infty} \mid a$); Σ_t is compressed for guidance ($Y = \pi^* \mid M_t$). Prediction is about *what will happen*; guidance is about *what to do*. Both are computed from the same history, with different targets.
- The information cost $I(\mathcal{C}_t; \Sigma_t)$ replaces $\text{DL}(\Sigma_t)$ as the theory-level compression term. The DL formulation remains useful as an *operational* cost measure for specific DAG encodings; it is not the theoretical quantity the IB objective minimizes.

Relationship between the two cost measures. Under MDL with a specific encoding scheme for DAGs (the one in 9.9), $DL(\Sigma_t)$ is an upper bound on $I(\mathcal{C}_t; \Sigma_t)$ for DAGs produced by the given encoder — coding cost dominates distinguishability cost. In practice, DL is computable and I is not, so the operational minimization uses DL as a proxy; the theoretical minimization uses I . The IB objective above is the theoretical statement; the DL-based minimization in 9.9 remains the practical one.

Variational form. The Shannon mutual information $I(\Sigma_t; \pi^* \mid M_t)$ in the relevance term collapses to zero when π^* is deterministic-from- M_t . The variational form (cf. 9.9) replaces the relevance term with the KL-divergence $D_{KL}(\pi^*(\cdot \mid M_t) \parallel Q_{\Sigma_t}(\pi \mid M_t))$ — note the π^* -first direction, which is *forced* by a regret-bound derivation (full derivation and admissible-divergence family analysis in T). Under bounded value range and deterministic π^* , Pinsker’s inequality gives $R(Q_{\Sigma_t}) \leq V_{\max} \sqrt{\frac{1}{2} D_{KL}(\pi^* \parallel Q_{\Sigma_t})}$ where R is the strategy-induced regret against π^* ; the opposite KL direction is vacuous ($+\infty$ under deterministic π^* whenever Q_{Σ_t} has off-optimum mass). The π^* -first KL is well-defined and graded under deterministic π^* . Under the variational reading, the AAT Σ_t is a tractable approximation of the policy-relevant posterior, and the KL term measures approximation quality — aligning the strategy compression with the variational free energy decomposition $-F = \text{accuracy} - \text{complexity}$ in active inference (Friston, FitzGerald, Rigoli, Schwartenbeck & Pezzulo 2017; Da Costa, Parr, Sajid, Veselic, Neacsu & Friston 2020). The direction alignment is convergent: both AAT’s regret-bound and active inference’s variational-free-energy derivations pick π^* -first KL (reverse-KL in the variational-inference vocabulary), with AAT’s additional interpretation being an upper-regret-bound rather than free-energy-gradient. The shared-IB-shape framing of the four AAT compression operations is the rate-distortion specialization of this variational picture; AAT’s commitment is to the rate-distortion form (which gives the (P1) Lagrangian-dual derivation and the four-instance unification at U-medium), not to the full variational free-energy interpretation.

Composition admissibility (P1) as IB Lagrangian-dual.

13.1 condition (P1) is currently stated as a lower-bound constraint:

Derived (p1-ib-dual, from composition-closure + rate-distortion duality).

$$I(\Lambda_x(X_{\text{micro},t}; \Lambda_o(o_{\text{micro},t+1}) \mid \Lambda_a(a_{\text{micro},t})) \geq (1 - \epsilon_I) \cdot I(X_{\text{micro},t}; o_{\text{micro},t+1} \mid a_{\text{micro},t}) \quad (\text{AA.4})$$

This is the *constraint form* of an IB problem. In Lagrangian form:

$$\Lambda^* \in \arg \min_{\Lambda \in \mathcal{P}} [I(X_{\text{micro}}; \Lambda_x(X_{\text{micro}})) - \beta(\epsilon_I) \cdot I(\Lambda_x(X_{\text{micro}}); Y_{\text{rel}})] \quad (\text{AA.5})$$

where $Y_{\text{rel}} = (o_{\text{micro},t+1} \mid a_{\text{micro},t})$ and $\beta(\epsilon_I)$ is the Lagrange multiplier corresponding to the relevance-preservation tolerance ϵ_I . The correspondence $\epsilon_I \leftrightarrow \beta$ is the standard rate-distortion duality: smaller ϵ_I (more relevance preserved) corresponds to larger β (less aggressive compression).

Consequence: admissible projections are those that sit *on or above* the IB frontier at rate $I(X; T) \leq I_{\max}(\epsilon_I)$. The information-theoretic content of (P1) is exactly “project onto the IB frontier with a tolerance ϵ_I .” This formalizes the connection previously logged in 13.1’s Working Notes and in 14.1’s §Connection to the Information Bottleneck.

What this resolves. The 13.1 Working Note “Open: Information Bottleneck unification” is now resolved for (P1): it is the Lagrangian-dual of the IB constraint at $\beta(\epsilon_I)$. The corresponding Working Note in 14.1 §6 moves from conjecture to derived result. Nothing else about (P1) changes — the condition continues to define admissible projections; only its information-theoretic reading is now explicit.

The symmetric two-Kalman witness, and the falsification handle.

The (P1)↔IB reading has a worked witness in the two-Kalman composition setup (14.1's linear-Gaussian case). With symmetric process noise $Q = q \begin{pmatrix} 1 & \rho \\ \rho & 1 \end{pmatrix}$, identical observation noise, and homogeneous gains ($\Delta K = 0$), the canonical-correlation basis between the micro-state $X = (\hat{\omega}_1, \hat{\omega}_2)$ and the relevance variable $Y = o_{t+1}$ is the symmetric pair $u_{\pm} = (1, \pm 1)/\sqrt{2}$, with generalized eigenvalues

Derived (linear-Gaussian-symmetric means-sum optimality).

$$\lambda_+ = \frac{(1 + \rho_M)\sigma_{\hat{\omega}}^2 + q(1 + \rho) + r}{q(1 + \rho) + r}, \quad \lambda_- = \frac{(1 - \rho_M)\sigma_{\hat{\omega}}^2 + q(1 - \rho) + r}{q(1 - \rho) + r}, \quad (\text{AA.6})$$

so the means-sum projection $u_+^\top X = (\hat{\omega}_1 + \hat{\omega}_2)/\sqrt{2}$ — the natural (P1)-admissible projection — is simultaneously the PCA-optimal 1D projection of X and the CCA-optimal 1D projection of X for predicting Y . The coincidence is symmetry-forced (the canonical basis and the principal-component basis of Σ_X align under symmetric Q and $\Delta K = 0$); in general the two bases differ. The retained relevance at 1D is $\frac{1}{2} \log \lambda_+$ of the full $\frac{1}{2} \log(\lambda_+ \lambda_-)$, ranging from half (at $\rho \rightarrow 0$, symmetric contributions) to all of it (at $\rho \rightarrow 1$, collapsed directions). Whether the means-sum additionally attains the Gaussian IB *frontier* remains the deferred mechanical derivation (Working Notes below); the canonical-direction identity here is the structural half that does not wait on it.

In the heterogeneous case ($\Delta K \neq 0$, where the means-sum produces genuine closure defect $\varepsilon_x > 0$), the symmetry breaks and the means-sum is *no longer* the top canonical direction: the IB-optimal 1D projection is a rotation of it, and — if the IB reading of (P1) is right — that rotation achieves strictly lower ε_x than the means-sum. This is a falsifiable consequence of the unification, checkable numerically in the existing two-Kalman machinery: if the rotated projection fails to improve ε_x when $\Delta K \neq 0$, the IB reading of (P1) is disconfirmed as more than shape-coincidence. Under the unification, means-sum admissibility is *conditional on gain symmetry* rather than free.

Hypothesis (heterogeneous-gain falsification handle).

What stays separate from the IB frame. Three admissibility and structural conditions do *not* reduce to IB:

- **(P2) Lipschitz continuity.** Not an IB constraint. The bridge lemma in 13.1 requires (P2) for analytic reasons (propagating bounded closure defect into bounded trajectory error); IB does not impose any continuity condition on compressors. (P2) remains a separate admissibility condition.
- **(P3) Dimensional reduction.** In the Gaussian-IB case relevant to composition, the IB-optimal T at any finite β typically uses full support of $\mathbb{R}^{\dim X}$; the categorical dimensionality reduction $\dim \mathcal{X}_c < \dim \mathcal{X}_{\text{micro}}$ is a *harder* condition than any rate constraint. (P3) remains separate. (In discrete cases it may be rate-implied, but the composition instance is Gaussian.)
- **Interventional relevance (Level 2).** The relevance variable in all four instances is associational (Y in a joint distribution with X and T). Strategy edges in the regime-indexed interpretation (9.5) want interventional relevance for Regime A: “what edge (i, j) predicts *under* $do(i)$.” This is a strictly stronger requirement than IB provides. Adapting IB to interventional relevance (causal IB, Wieczorek & Roth 2017 and follow-ups) is an extension direction, not a specialization of the master IB.

The honest slogan is therefore “the (P1)-analog in each compression operation is IB; regularity, dimensionality, and interventional relevance are separate conditions that compose with it.”

Epistemic Status.

Robust-qualitative. The claim that the four compression operations share IB shape is *discussion-grade* — it is a presentational observation supported by examining each segment’s formulation and noting structural alignment. The Σ_t source reformulation is a *formulation choice* replacing one underspecified source with a cleaner one; no derivation is needed because the new formulation and the old are not the same claim. The (P1) as Lagrangian-dual of IB is *derived* — rate-distortion duality is standard (see §1.12–13 of Cover & Thomas) and the constraint-form $\leftrightarrow$ Lagrangian-form equivalence is mechanical.

Max attainable: *robust-qualitative* for the shared-shape claim; *derived conditional on rate-distortion duality* for (P1) as IB Lagrangian-dual; *formulation* for the Σ_t source reformulation. The strongest version of the unification claim — “U-strong,” that the four operations are the same optimization problem with different bindings — is *not* established and is unlikely to be, for reasons in the Discussion.

This segment does not promote any of the four instance segments. Each remains the primary site for its own instance. What this segment adds is the shared-shape framing, the Σ_t source fix, and the (P1) derivation.

Discussion.

Why not go further. Two distinct claims could be made: (a) U-strong — the four operations are the *same* optimization problem with different bindings, so a single derivation specializes to each instance; (b) U-medium — the four share *shape*, vocabulary, and variational structure but differ in substantive ways (source type, relevance-variable availability, computability) that prevent a single derivation from covering all four. U-strong requires cross-instance theorems, results of the form “because these are all IB, property X holds across all four.” No such theorems are identified. The shared IB shape is legibility-enhancing but does not produce deductions for free. This segment therefore states U-medium (shared shape) honestly rather than overclaiming U-strong.

Honest credit to the hierarchical-generative-model lineage. AAT’s four compression operations are a structurally narrower family than the operations a hierarchical generative model in the predictive-coding lineage (Friston 2008, “Hierarchical models in the brain,” *PLoS Comp. Biol.* 4; Friston 2010, “The free-energy principle: a unified brain theory?” *Nat. Rev. Neurosci.* 11; Clark 2013, “Whatever next?,” *Behav. Brain Sci.* 36; Hohwy 2013, *The Predictive Mind*, OUP) could express. A hierarchical generative model layers compressions of compressions, with each layer producing a representation tuned to the next layer’s prediction target — and the AAT compressions (M_t for prediction, Σ_t for guidance, shared intent for coordination, Λ for level-bridging) are all expressible within that frame as specific layer-bindings. What AAT adds is structural: (a) the *relevance variables* Y are made first-class with explicit per-instance bindings (see the table in §“The shared IB shape”); (b) the (P1)–(P3) admissibility conditions for composition give a measurable closure-defect bound that hierarchical-generative-model layering does not natively produce (13.1); (c) regime-indexed edges (9.5) introduce Pearl-Level-2 relevance for Regime A, which standard hierarchical generative models do not address. The shared family is real; AAT’s additions are also real. The honest framing is “AAT’s compressions are a structured subset of the hierarchical-generative-model family with additional structure load-bearing for AAT’s specific results.”

Why this is not the “biggest available unification.” It is better described as the *most legible* unification. The pattern is already half-named at each instance site (three of four segments already write their objectives in IB form); the segment-level contribution is sharpening the naming and fixing one formulation issue (Σ_t source), not discovering a new structural result. The biggest leverage concision move in the theory — the sector-Lyapunov template in G — is substantively different: it removes repeated content and absorbs proofs that previously appeared in multiple segments. This segment adds legibility without removing much.

What each β means. The four trade-off parameters are not unified in value, only in role:

- $\beta(\rho, \pi)$ for M_t depends on environmental volatility and policy; fast-changing environments or high-value policies both warrant higher β .

- β_Σ for Σ_t is the cognitive cost per decision-relevant bit retained; bounded-context agents (LLMs with finite windows) have low β_Σ .
- The bandwidth β for shared intent is the communication-capacity-per-coordination-bit; low-bandwidth teams have low β .
- The $\beta(\epsilon_I)$ for (P1) is the rate-distortion Lagrange multiplier at tolerance ϵ_I .

These are four calibration problems, not one. A theory of *resource allocation across the four compression operations* — where a single cognitive budget is split across modeling, strategy, communication, and composition — is not currently in AAT, and the unification does not produce it. If such a theory is ever needed, it would tie the four β s together through an outer optimization; the spike flags this as an open direction.

Relationship to 14 and 14.1. The unity dimensions parametrize rate-distortion curves for closure defects (14.1). Under the present segment’s framing, unity dimensions are *curve parameters for the Λ instance of the IB schema*. This tightens the existing 14.1 Discussion, which already presented the rate-distortion reading as the correct interpretation; it is now stated uniformly across the four compression operations.

Relationship to Q. The DAG-uniqueness result derives the *shape* of Σ_t from operational postulates (Q, Cox-analog). The present segment characterizes the *compression level* of Σ_t . These operate at different levels: uniqueness supplies the representational type (the compressed T is a DAG), IB supplies the rate-distortion trade-off within that type (how richly the DAG is populated). The two arguments are independent and compatible — IB-compressed Σ_t is a DAG-with-fewer-nodes, not a non-DAG.

Position relative to the additive-coordinate-forcing pattern. The IB Lagrangian’s $I(X; T) - \beta \cdot I(T; Y)$ is an additive decomposition on a log-scale coordinate (mutual information is logarithmic by construction) — the same shape that 5.5 catalogs at the chain / divergence / update layers. The IB case is classified as an *adjacent family member* rather than a fourth primary instance because 2.2 adopts the Lagrangian form from Tishby, Pereira & Bialek 1999 as an applied external theorem rather than re-deriving it under an AAT-internally-motivated additivity axiom. The Shore-Johnson 1980 system-independence axioms that axiomatize IB’s additivity are cited in T §6.1; if a future framing pass promotes Shore-Johnson to an explicit AAT axiom, IB would move from adjacent member to fourth primary instance.

Working Notes.

- **The Σ_t source reformulation changes the theoretical quantity without changing the operational computation.** 9.9’s DL-based minimization remains the practical procedure; what changes is the theoretical claim about what that procedure approximates. Depending on promotion-workflow preferences, the reformulation may warrant a small note in 9.9 referencing this segment.
- **(P1) as IB Lagrangian-dual does not close 13.1’s other Working Notes.** Computability of (P1) for nonlinear/non-Gaussian systems, choice of ϵ_I , and N -agent scaling remain open. The unification resolves the *framing* question but not the *computation* questions.
- **Synthesis segment vs. full unification.** This segment is a targeted synthesis: it fixes the Σ_t source, derives (P1) as IB Lagrangian-dual, and tables the four instances. It does *not* rewrite the four instance segments as specializations of a master schema. If future work reveals that the four instance segments are pulling in inconsistent directions (e.g., if the Σ_t reformulation creates friction with the DL-based minimization in 9.9), that is the signal to escalate to a full rewrite in which the instance segments become thin specializations. Short of that signal, this targeted synthesis is the stopping point.
- **The Gaussian IB frontier derivation is deferred.** Completing the derivation that the means-sum projection attains the Gaussian IB frontier in the symmetric two-Kalman case would promote 14.1’s rate-distortion claim from discussion-grade to derived. The obstruction is mechanical (specializing Chechik, Globerson, Tishby & Weiss 2005, “Information Bottleneck for Gaussian Variables,” *J. Machine Learning Research* 6:165–188, Theorem 3.1 to the symmetric-gain setup) rather than structural — feasible as a follow-up derivation, not required for this segment. The *canonical-direction half* (means-sum =

PCA-optimal = CCA-optimal under symmetry, with explicit $u_{\pm}/\lambda_{\pm}$) is now landed in the body's §symmetric-two-Kalman-witness (2026-07-16, from the bulk-64 verification of spikes/.integrated/spike-ib-unification-plan.md §5.4–5.6); only the frontier-attainment step remains deferred.

- **The heterogeneous-gain numeric check is unrun.** The falsification handle in the body (§witness: IB-optimal rotation should beat the means-sum's ε_x when $\Delta K \neq 0$) is a small numerical experiment on the existing two-Kalman machinery — a natural companion to the frontier derivation above; neither has been executed.
-

Appendix AB

Derivation: Observation-Ambiguity Bias-Bound Constant C

The observation-ambiguity bias bound carried by Class 3 (Coupled) agents in the logogenic-agents scope — $\|\Delta M_{\text{bias}}\| \leq C \cdot \kappa_{\text{processing}} \cdot I(G; \Omega_{\tau} \mid e_{\tau}, M_{\tau-})$ (`#scope-observation-ambiguity-modulation`, `#result-section-ii-survival`) — previously treated the constant C as “domain-dependent” and the bound as “order-of-magnitude guidance, not a theorem.” This appendix derives C under two named sub-scopes, records a no-go showing that C cannot be universal without the (PI) parameterization-invariance axiom of 4.10, and documents two failed derivation routes so future agents do not repeat them.

Formal Expression.

§1 — Setup and the audit before strengthening. Let $M_{\tau-} \in \mathcal{M}$ be the pre-update epistemic substate. Let e_{τ} be the event triggering update, G_t the goal substate, and Ω_{τ} the latent world-state. At the **causally-disciplined** baseline ($\kappa_{\text{processing}} = 0$ — Class 1 by construction or an idealized Class 2 in effect, equally; 6.3), the update $f_M(M_{\tau-}, e_{\tau})$ is goal-blind: $M_{\tau+}^{\text{decoupled}} = f_M(M_{\tau-}, e_{\tau})$, and the bias below is zero. In Class 3 (Coupled) scope, the coupled update $f_X^M(X_{\tau-}, e_{\tau})$ carries goal-conditional reweighting; the bias is

$$\Delta M_{\text{bias}} := f_X^M(X_{\tau-}, e_{\tau}) - f_M(M_{\tau-}, e_{\tau}) \quad (\text{AB.1})$$

Pre-strengthening type audit. For the bound “ $\|\Delta M_{\text{bias}}\| \leq C \cdot \kappa \cdot I$ ” to be well-typed:

1. **Norm on $\mathcal{M}$.** The LHS is a norm on model-space. $\mathcal{M}$ is the model space of 2.1. Three candidate norms: Euclidean on parameters, total variation on induced measures, Fisher-Rao geodesic distance. **Under the (PI) parameterization-invariance axiom in 4.10** (fourth primary instance of 5.5), Euclidean-on-parameters is a coordinate artifact and Fisher-Rao is the canonical AAT-invariant choice on statistical-manifold sub-cases of $\mathcal{M}$ (Čencov 1982 uniqueness).
2. **Regularity of f_X^M .** For the bound to be a theorem, f_X^M must satisfy some regularity. The Bayesian-posterior model (prior + likelihood reweighting) is the canonical working class; it covers the attention-reweighting mechanism typical of Class 3 (Coupled) architectures to leading order.
3. **Information coordinate.** $I(G; \Omega_{\tau} \mid e_{\tau}, M_{\tau-})$ is in nats (or bits). Converting information to state-distance is exactly what the Pinsker / Bretagnolle-Huber / Otto-Villani / Bakry-Émery machinery does.

The bound is well-typed under (a) named norm, (b) bounded f_X^M regularity, (c) specified information-geometric inequality. Without (PI) adopted, C is norm-dependent; under (PI), Fisher-Rao is canonical and C becomes derivable.

§2 — Track 1: Transport-inequality derivation under log-Sobolev + Lipschitz-posterior.

Named sub-scope (H1-H3).

Derived (w2-transport-track, conditional on H1-H3).

- **(H1) Statistical-manifold sub-case.** Each $M_t \in \mathcal{M}$ corresponds to a probability distribution P_{M_t} over world-states. $\mathcal{M}$ is (locally) a statistical manifold.
- **(H2) Log-Sobolev inequality.** The observation distribution $P_{\Omega|e,M}$ satisfies a log-Sobolev inequality (LSI) with constant $\rho_{\text{LSI}} > 0$. Sufficient condition: $P_{\Omega|e,M}$ is strongly log-concave with constant K (Bakry-Émery 1985 curvature-dimension condition gives $\rho_{\text{LSI}} \geq K$). For Gaussian observation models, $\rho_{\text{LSI}} = 1/\sigma^2$ explicitly.
- **(H3) Lipschitz-posterior stability.** The Bayesian-posterior pushforward from observation to state is L_{post} -Lipschitz in W_2 :

$$W_2(P_{M|e,G}, P_{M|e}) \leq L_{\text{post}} \cdot W_2(P_{\Omega|e,G}, P_{\Omega|e}) \quad (\text{AB.2})$$

Sufficient condition: well-posed Bayesian inverse problem with bounded log-likelihood Hessian (Stuart 2010 *Inverse problems: a Bayesian perspective*, *Acta Numerica* 19:451–559, Theorem 4.6; Hairer, Stuart & Vollmer 2014 *SIAM J. Math. Anal.* 46(1):415–451 for explicit W_2 -bounds in infinite-dimensional settings).

Step 1 — KL from mutual information. By the chain rule of relative entropy (Cover & Thomas 2006 Theorem 2.5.3):

$$\mathbb{E}_G[\text{KL}(P_{\Omega|e,M,G} \| P_{\Omega|e,M})] = I(G; \Omega_\tau | e_\tau, M_{\tau-}) \quad (\text{AB.3})$$

Exact identity.

Step 2 — Otto-Villani under LSI. From (H2) and Otto & Villani 2000 *J. Funct. Anal.* 173(2):361–400 Theorem 1:

$$W_2^2(P_{\Omega|e,M,G}, P_{\Omega|e,M}) \leq \frac{2}{\rho_{\text{LSI}}} \cdot \text{KL}(P_{\Omega|e,M,G} \| P_{\Omega|e,M}) \quad (\text{AB.4})$$

Taking expectation over G and substituting Step 1:

$$\mathbb{E}_G[W_2^2(P_{\Omega|e,M,G}, P_{\Omega|e,M})] \leq \frac{2}{\rho_{\text{LSI}}} \cdot I(G; \Omega_\tau | e_\tau, M_{\tau-}) \quad (\text{AB.5})$$

Step 3 — Lipschitz-posterior pushforward. From (H3):

$$\mathbb{E}_G[W_2^2(P_{M|e,G}, P_{M|e})] \leq L_{\text{post}}^2 \cdot \mathbb{E}_G[W_2^2(P_{\Omega|e,G}, P_{\Omega|e})] \quad (\text{AB.6})$$

Step 4 — $\kappa_{\text{processing}}$ factor. The $\kappa_{\text{processing}}$ coefficient from 6.3's Class 1/2/3 (Separated/Partial/Coupled) taxonomy multiplies the goal-conditional reweighting strength by the modularity coefficient. For Class 3 (Coupled), $\kappa_{\text{processing}} \approx 1$; for Class 2 (Partial) with modularity κ , the effective bound carries a factor κ multiplicatively on the information term.

Result (Track 1). [Derived, conditional on (H1)-(H3)]

$$\mathbb{E}[W_2^2(M_{\tau+}^{\text{coupled}}, M_{\tau+}^{\text{decoupled}})] \leq \frac{2L_{\text{post}}^2}{\rho_{\text{LSI}}} \cdot \kappa_{\text{processing}} \cdot I(G; \Omega_{\tau} \mid e_{\tau}, M_{\tau-}) \quad (\text{AB.7})$$

The constant is explicit: $C_{W_2}^2 = 2L_{\text{post}}^2/\rho_{\text{LSI}}$, linear in I , with geometric interpretation:

- L_{post} grows with *prior-likelihood tension* — ill-conditioned inverse problems amplify bias.
- ρ_{LSI} grows with *observation-distribution concentration* — sharper observations yield tighter bounds.
- The ratio $2L_{\text{post}}^2/\rho_{\text{LSI}}$ captures the **geometric stiffness** of the update: small- ρ /large- L = soft, goal-sensitive updates; large- ρ /small- L = stiff updates that resist goal-bias.

§3 — Track 2: Fisher-Rao derivation under (PI) + Čencov + small-information regime.

Named sub-scope (H1 + H4).

Derived (fisher-rao-track, conditional on H1 + H4).

- **(H1) Statistical-manifold sub-case** — as in §2. Under the (PI) parameterization-invariance axiom of 4.10, Čencov’s 1982 uniqueness theorem (*Statistical Decision Rules and Optimal Inference*, AMS) forces the Fisher information metric as the canonical Riemannian metric on $\mathcal{M}$ up to global scale (Ay, Jost, Lê & Schwachhöfer 2017 *Information Geometry*, Theorem 5.1).
- **(H4) Small-information regime.** $I(G; \Omega_{\tau} \mid e_{\tau}, M_{\tau-}) \ll 1$ nat. The second-order Taylor expansion of KL at coincident distributions is sharp.

Step 1 — KL-to-Fisher-squared-distance identity. For nearby distributions P, Q on a statistical manifold with Fisher metric $\mathbf{I}$, the KL divergence admits the second-order expansion

$$\text{KL}(P\|Q) = \frac{1}{2} \cdot d_{FR}^2(P, Q) + O(d_{FR}^3) \quad (\text{AB.8})$$

where d_{FR} is the Fisher-Rao geodesic distance (Cover & Thomas 2006 §12.5; Amari & Nagaoka 2000 §3.7 Theorem 3.1). This is the infinitesimal form of the Bregman divergence on the exponential family’s dual geometry (cf. T §6.3 for the related Fenchel-Bregman identification).

Step 2 — Posterior-displacement in Fisher-Rao. Under (H1)+(H4) + Bayesian-posterior-class f_X^M , the goal-conditional posterior and goal-marginalized posterior are nearby points on the statistical manifold; their Fisher-Rao distance satisfies

$$d_{FR}^2(M_{\tau+}^{\text{coupled}}, M_{\tau+}^{\text{decoupled}}) \leq 2 \cdot \text{KL}(P_{M|e,G} \parallel P_{M|e}) + O(d_{FR}^3) \quad (\text{AB.9})$$

Using the data-processing inequality $\text{KL}(P_{M|e,G} \parallel P_{M|e}) \leq \text{KL}(P_{\Omega|e,G} \parallel P_{\Omega|e})$ (Bayesian posterior is a pushforward), and Step 1 of §2 ($\mathbb{E}_G[\text{KL}(P_{\Omega|e,M,G} \parallel P_{\Omega|e,M})] = I$):

$$\mathbb{E}[d_{FR}^2(M_{\tau+}^{\text{coupled}}, M_{\tau+}^{\text{decoupled}})] \leq 2 \cdot I(G; \Omega_{\tau} \mid e_{\tau}, M_{\tau-}) \cdot (1 + o(1)) \quad (\text{AB.10})$$

Step 3 — Taking square roots (with κ factor). Under Class 2 (Partial) modularity, $\kappa_{\text{processing}}$ enters multiplicatively on the goal-conditional reweighting strength, hence on the KL and Fisher-Rao distance:

Result (Track 2). *[Derived, conditional on (H1)+(H4), small-information regime]*

$$\mathbb{E} \|\Delta M_{\text{bias}}\|_{FR} \leq \sqrt{2} \cdot \sqrt{\kappa_{\text{processing}} \cdot I(G; \Omega_{\tau} | e_{\tau}, M_{\tau-}) \cdot (1 + o(1))} \quad (\text{AB.11})$$

The constant is universal and dimension-free: $C_{FR} = \sqrt{2}$ (nats) or $\sqrt{2 \ln 2}$ (bits). **No domain-specific parameters.** The (PI) commitment eliminates the coordinate-dependence that made C ambiguous under Euclidean-parameter norms.

Scaling difference between tracks. Track 1 gives $W_2^2 \propto I$ (linear in information). Track 2 gives $d_{FR}^2 \propto I$, equivalently $d_{FR} \propto \sqrt{I}$ (square-root in information). The two tracks are not contradictory — they are bounds under different assumptions in different metrics. Track 1 is tight in the large- I regime where Otto-Villani’s linear form dominates; Track 2 is tight in the small- I regime where the Fisher-metric second-order expansion is sharp. The two compose at intermediate scales by taking the tighter of the two under the local curvature conditions.

§4 — No-go result: universal C under Euclidean-parameter norm fails.

Claim. No universal constant C — independent of $\mathcal{M}$ ’s geometry, parameterization, and coupled-update structure — exists such that $\|\Delta M_{\text{bias}}\|_{\text{Euclidean}} \leq C \cdot \kappa \cdot I$ under the Euclidean norm on an arbitrary parameter vector.

Proved (no-universal- C -euclidean, counterexample-grade).

Counterexample. Let $\mathcal{M} = \{N(0, \sigma^2) : \sigma > 0\}$ parameterized by σ . The Fisher information in σ scales as $\mathbf{I}(\sigma) = 2/\sigma^2$ — small near large σ , large near small σ . For any $C_0 < \infty$, choose σ_0 large enough that the Fisher-Rao displacement corresponding to $I(G; \Omega) = 1$ nat of goal-conditional reweighting translates to Euclidean- σ displacement $\sigma_0/\sqrt{2} > C_0$. Specifically, by Track 2’s $\sqrt{2}$ bound:

$$d_{FR}(M^{\text{coupled}}, M^{\text{decoupled}}) \leq \sqrt{2} \cdot 1 = \sqrt{2} \quad (\text{AB.12})$$

Under the Fisher metric $\mathbf{I}(\sigma) = 2/\sigma^2$, the Fisher-Rao line element is $ds^2 = 2d\sigma^2/\sigma^2$, so $\Delta\sigma \approx \sigma \cdot d_{FR}/\sqrt{2} = \sigma$ (Euclidean displacement scales linearly in σ). Taking $\sigma \rightarrow \infty$ gives arbitrarily large Euclidean- σ displacement for fixed $I = 1$ nat. No universal C in Euclidean-parameter norm exists.

Implication. The no-go strengthens the (PI) commitment rather than weakening the bound: **under (PI), Fisher-Rao is the canonical norm and $C = \sqrt{2}$ is universal; without (PI), C does not exist as a universal constant.** The (PI) axiom is *load-bearing* for this derivation, not coincidental.

Position relative to 5.1. The no-go has the shape of a floor-pattern no-go (external obstruction: Euclidean-parameter norms carry unbounded Fisher condition numbers; escape: (PI) + Fisher-Rao gives universal C), but it does not match the five-element test for a floor instance (see 5.1): it has a *single* escape (the (PI) adoption), not ≥ 2 distinct escapes, and its strengthened consequence is *re-use* of existing (PI)+Čencov machinery (fourth primary instance of 5.5) rather than elevation of new machinery. The honest position: this no-go is a **downstream theorem of the (PI) commitment**, not a new floor instance. It belongs in this appendix as motivating justification for why (PI) is load-bearing for the derivation, not as a separate meta-segment entry. (The floor-vs-coordinate-forcing distinction is articulated explicitly in #disc-additive-coordinate-forcing §”Downstream consequences of the (PI) commitment”.)

§5 — Failed derivation routes (honest record). Two derivation routes were attempted and failed at a structural level. Documented here so future agents do not re-attempt them without new evidence.

(F1) Cramér-Rao inversion — wrong direction. Attempt: treat G as a parameter estimated from Ω given (e, M) ; use the Cramér-Rao lower bound $\text{Var}(\hat{G}) \geq \mathbf{I}_G^{-1}$ inverted to bound $\|\Delta M_{\text{bias}}\| \leq L_M \cdot \sqrt{\mathbf{I}_G^{-1}}$ via a posterior-sensitivity factor. Failure: Cramér-Rao bounds estimator error *below*, not above; the bound goes in the opposite direction to what the bias bound needs. Inverting Cramér-Rao for an upper bound on bias propagation requires a fixed-estimator-class assumption that is not available for a Class 3 (Coupled) agent’s attention-reweighting mechanism. The transport-inequality route of §2 is the correct machinery.

(F2) Rate-distortion inversion — wrong problem structure. Attempt: apply Shannon’s rate-distortion inequality $R(D) \geq h(X) - \frac{1}{2} \log(2\pi e D)$ (Gaussian source), invert to $D \leq \sigma^2 \cdot 2^{-2(R-h(X))}$ with “bits in” = $I(G; \Omega \mid e, M)$. Failure: rate-distortion theory describes the minimum bits-per-symbol required to represent a source within distortion D ; it does **not** describe the maximum distortion induced by injecting side-information into an update. Source-coding theorems are about optimal representations of a source, not spatial displacement induced by side-information injection. The direction is again wrong. Transport inequalities (Attempt B) are the correct machinery.

The pattern shared across (F1) and (F2): **information-theoretic source-coding theorems (Cramér-Rao bounds, rate-distortion lower bounds) are structural lower-bound machinery for an estimator’s or encoder’s performance; they cannot be inverted to produce upper bounds on displacement induced by side-information.** Upper bounds on displacement require transport-inequality machinery (Pinsker, Otto-Villani, Bakry-Émery, posterior-stability).

§6 — Gaussian worked example.

Setup. Gaussian observation model $P_{\Omega|e,M} = N(\mu_M, \sigma^2 I)$ with fixed σ^2 ; conjugate-Gaussian prior $P_M = N(0, \tau^2 I)$; goal-conditional likelihood reweighting $P_{\Omega|e,M,G} = N(\mu_M + \beta(G), \sigma^2 I)$ where $\beta(G)$ is a goal-dependent shift with $\|\beta(G)\|^2$ bounded by the information budget $I(G; \Omega_\tau)$.

Illustrative, exact under specified assumptions.

LSI constant. Gaussian $N(\mu, \sigma^2 I)$ satisfies LSI with $\rho_{\text{LSI}} = 1/\sigma^2$ (Bakry-Émery; direct Hessian bound on $-\log p$).

Posterior-Lipschitz constant. Conjugate-Gaussian posterior: $P_{M|\Omega} = N(\mu_{\text{post}}, \Sigma_{\text{post}})$ with $\mu_{\text{post}} = \tau^2/(\tau^2 + \sigma^2) \cdot \Omega$. Lipschitz-in- Ω constant $L_{\text{post}} = \tau^2/(\tau^2 + \sigma^2) < 1$. Under W_2 -Lipschitz (Hairer-Stuart-Vollmer 2014), $L_{\text{post}} \leq 1$ for well-conditioned priors.

Track 1 bound (numerical). $C_{W_2}^2 = 2L_{\text{post}}^2/\rho_{\text{LSI}} = 2\sigma^2 \cdot \tau^4/(\tau^2 + \sigma^2)^2$. In the prior-dominant limit ($\tau^2 \ll \sigma^2$): $C_{W_2}^2 \approx 2\tau^4/\sigma^2$ (tight prior amplifies bias). In the likelihood-dominant limit ($\sigma^2 \ll \tau^2$): $C_{W_2}^2 \approx 2\sigma^2$ (sharp observations limit bias).

Track 2 bound (numerical). Under (PI) + Fisher-Rao on the Gaussian mean-manifold with metric $\mathbf{I}(\mu) = I/\sigma^2$: $d_{FR}(\mu_1, \mu_2) = \|\mu_1 - \mu_2\|/\sigma$. Under small- I : $\mathbb{E} d_{FR}(\Delta M_{\text{bias}}) \leq \sqrt{2 \cdot \kappa \cdot I}$. In Euclidean-on- μ norm: $\|\Delta \mu_{\text{bias}}\| \leq \sigma \cdot \sqrt{2\kappa I}$ — the σ prefactor is the Fisher-Rao-to-Euclidean conversion. The Euclidean form is coordinate-dependent (per Attempt E §4); Fisher-Rao form is universal.

Comparison at operating point. For $\tau = \sigma = 1$ (balanced): Track 1 gives $C_{W_2}^2 = 2 \cdot (1/4) = 1/2$, so $\mathbb{E} W_2 \leq \sqrt{I/2}$. Track 2 gives $\mathbb{E} d_{FR} \leq \sqrt{2I}$. The two are within a factor of 2 of each other at balanced scale; they diverge in the ill-conditioned limits where Track 1’s L_{post} or ρ_{LSI} blows up while Track 2 remains dimension-free.

What Is Derived vs. What Is Chosen.

Table AB.1

Property	Source	Strength
Chain-rule identity $\mathbb{E}_G[\text{KL}(P_{\Omega G} P_{\Omega})] = I(G; \Omega \cdot)$	Cover & Thomas 2006 Theorem 2.5.3	Exact
Pinsker's inequality $\ P - Q\ _{TV} \leq \sqrt{\frac{1}{2}\text{KL}}$	Standard (Csiszár-Körner)	Exact
Otto-Villani $W_2^2 \leq (2/\rho_{\text{LSI}})\text{KL}$ under LSI	Otto & Villani 2000 Theorem 1; Bakry-Émery 1985 curvature-dimension for LSI	Exact
KL-to-Fisher-squared-distance $\text{KL}(P\ Q) = \frac{1}{2}d_{FR}^2 + O(d_{FR}^3)$	Cover-Thomas 2006 §12.5; Amari-Nagaoka 2000 §3.7 Theorem 3.1	Exact (standard information geometry)
Bayesian-posterior W_2 -Lipschitz stability L_{post}	Stuart 2010 Theorem 4.6 under well-posed inverse problem	Conditional on (H3)
Track 1 bound $\mathbb{E}[W_2^2(M^{\text{coupled}}, M^{\text{decoupled}})] \leq (2L_{\text{post}}^2/\rho_{\text{LSI}}) \cdot \kappa \cdot I$	Composition of §2 Steps 1–4	Derived (conditional on H1-H3)
Track 2 bound $\mathbb{E} \ \Delta M_{\text{bias}}\ _{FR} \leq \sqrt{2} \cdot \sqrt{\kappa I}(1 + o(1))$	Composition of §3 Steps 1–3 under (PI) + Čencov	Derived (conditional on H1+H4, small-I regime)
$C_{FR} = \sqrt{2}$ is universal, dimension-free, no domain constants	§3 Result	Exact under (H1)+(H4)
$C_{W_2}^2 = 2L_{\text{post}}^2/\rho_{\text{LSI}}$ with explicit geometric interpretation	§2 Result	Exact under (H1)-(H3)
Attempt E no-go: no universal C in Euclidean-parameter norm	Heteroscedastic-normal counterexample §4	Proved (counterexample-grade)
(F1) Cramér-Rao inversion fails	Direction mismatch (estimator lower bound cannot invert to upper bound)	Recorded failure
(F2) Rate-distortion inversion fails	Problem-structure mismatch (source-coding theorem cannot yield side-information injection bound)	Recorded failure
Gaussian worked example with explicit C_{W_2} and C_{FR}	Conjugate-Gaussian direct computation	Exact under specified assumptions

The dividing line: both tracks are derived theorems under named hypotheses; the hypotheses (LSI, Lipschitz-posterior, small-I, statistical-manifold sub-case, (PI) adoption) are either standard mathematical regularity (LSI, Lipschitz-posterior) or AAT-internal axioms ((PI), already adopted in 4.10 and elevated to fourth primary instance of 5.5). The no-go §4 justifies the (PI) commitment as load-bearing rather than coincidental. The failed attempts §5 document structural reasons two alternative routes cannot work, preventing future re-attempts.

Epistemic Status.

Conditional — exact under H1-H3 (Track 1) or H1+H4 (Track 2). Both tracks are derived theorems under named hypotheses:

- Track 1 is exact under the three regularity conditions (H1 statistical-manifold sub-case; H2 log-Sobolev on observation distribution; H3 Lipschitz-posterior stability). All three are standard mathematical regularity conditions satisfied for canonical cases (Gaussian observations: (H2) direct; well-posed Bayesian inverse problems: (H3) per Stuart 2010). The bound is linear in I , with explicit geometric-stiffness constants.
- Track 2 is exact under (H1) + (H4) small-information regime, leveraging (PI) + Čencov. The bound is $\sqrt{I}$ scaling with universal dimension-free constant $C_{FR} = \sqrt{2}$. Does not depend on domain-specific L_{post} or ρ_{LSI} ; the price is the small-I regime restriction.

The no-go §4 is *proved* via the heteroscedastic-normal family counterexample — universal C in Euclidean-parameter norm does not exist; the (PI) commitment is load-bearing for the derivation. The failed attempts §5 are *recorded* at structural-reason level; they document why Cramér-Rao inversion and rate-distortion inversion cannot produce the needed upper bound, not merely that the attempts failed.

Upgrade over prior formulation. The bias bound in `#result-section-ii-survival` and `#scope-observation-ambiguity-modulation` previously carried the honest admission “*Without the constant, the bound is order-of-magnitude guidance, not a theorem.*” This appendix **promotes the bound to theorem-level** under named sub-scopes. The order-of-magnitude admission is superseded by the two explicit tracks with their H-conditions. Updates to the logogenic-agents segments cross-reference this appendix.

Max attainable: *exact* for the composition of the three textbook inequalities (Step 1 chain-rule; Step 2 Otto-Villani; Step 3 Lipschitz-posterior or KL-to-Fisher-Rao identity) and for the no-go counterexample §4. *Conditional* overall — the theorem statements carry their H-conditions explicitly, and the sub-scope boundary is where “derived” transitions back to “admitted empirical regularity” (cases where H2 or H3 fails require separate treatment, flagged in Working Notes).

Honest failure modes.

The appendix does **not** cover:

1. **Non-statistical-manifold sub-cases of $\mathcal{M}$.** If $\mathcal{M}$ carries structured state-space geometry (mixed continuous/discrete components, graphical-model structure, non-parametric models without sufficient statistics), (H1) is weakened and Čencov’s uniqueness does not straightforwardly apply. The Fisher-Rao track degrades; Track 1 may still apply under suitable alternative regularity.
2. **Unbounded KL / structural-failure limit.** If $P_{\Omega|e,M,G}$ has support disjoint from $P_{\Omega|e,M}$ at G -values carrying non-trivial prior probability, the KL (and thus I) diverges. The bound is vacuous in this regime. This is a structural failure of the coupled-update model, not a bound weakness.
3. **Non-Lipschitz pushforward.** If f_X^M is not a well-posed Bayesian posterior (e.g., ill-posed inverse problem, degenerate likelihood, rank-deficient observation channel), (H3) fails and Track 1 does not apply. Track 2’s small- I expansion may still apply in the local-linearization sub-region.
4. **Large-information regime beyond the small- I expansion.** Track 2’s $\sqrt{I}$ scaling uses the second-order Fisher-metric expansion of KL; the full global relationship is lattice-dependent. A fully-global Fisher-Rao statement requires either compact-manifold assumption (Fisher-Rao distance bounded) or a staged-application argument (many small- I steps composed).
5. **Heavy-tailed observation distributions.** LSI fails for heavy-tailed measures. Track 1 degrades to the Pinsker bound $W_2 \leq \sqrt{KL/(2\rho_{\text{LSI}})}$ when $\rho_{\text{LSI}} \rightarrow 0$; this is the $O(\sqrt{\varepsilon})$ baseline without the LSI tightening.

Discussion.

Position in the 5.5 meta-pattern. Track 2’s $C_{FR} = \sqrt{2}$ result is a **downstream theorem of the fourth primary instance** (Čencov-invariance at the metric layer, under (PI)). Once (PI) adoption forces Fisher-Rao as the canonical metric on statistical-manifold sub-cases of $\mathcal{M}$, the KL-to-Fisher-squared-distance identity supplies the universal constant “for free” in the small- I regime. The no-go §4 tightens the connection: the (PI) commitment is not just “useful axiom” but load-bearing for the bound’s theorem-level status — without (PI), C does not exist as a universal constant.

Connection to AJ. Both segments use Pinsker’s inequality as a first step in propagating a KL bound. After Pinsker, the two segments diverge: this segment’s cascade (Pinsker $\rightarrow$ Otto-Villani $\rightarrow$ Lipschitz-posterior $\rightarrow W_2$ on pushforward) is a KL-to-state-distance machinery; `#deriv-variational-sector-condition`’s use (Pinsker + Cauchy-Schwarz $\rightarrow$ scalar

sector-constant degradation) is a KL-to-scalar-constant machinery. The shared first step is Pinsker; the post-Pinsker cascade is not shared, so `#deriv-variational-sector-condition` is an adjacent family member rather than a candidate for a shared `#posterior-displacement-template` (a possible future extraction recorded in Working Notes).

The $\varepsilon(\kappa)$ degradation for causally-disciplined results. This appendix *is* the degradation bound that the causal-discipline scoping rule of 6.3 calls for: a Part II result that holds exactly at $\kappa_{\text{processing}} = 0$ (for any causally-disciplined agent) degrades off that zero at a rate this bound controls — $\|\Delta M_{\text{bias}}\|$ enters multiplied by $\kappa_{\text{processing}}$, exact (zero bias) at the causally-disciplined value and decaying controllably as coupling grows. The first realized client is the value object’s Q_O G_I -independence (6.6): the Q_O degradation at $\kappa_{\text{processing}} > 0$ is the ΔM_{bias} this appendix bounds, propagated through the value functional. That a successful $\varepsilon(\kappa)$ bound exists is itself the evidence that the result-at-the-noun (causal discipline, the shared zero) is the right object — the result degrades continuously off the zero rather than requiring the architecture certificate.

Forward-looking clients. Three open extensions in 5.1 Working Notes are candidate future primary clients of this appendix’s machinery: causal-IB (KL $\rightarrow W_2$ on post-intervention pushforward), misspecification-cost (KL $\rightarrow W_2$ on posterior under model misspecification), composition-scope-robustness (KL $\rightarrow W_2$ on composite-state pushforward). Each would reuse the Track 1 Otto-Villani + Lipschitz-posterior cascade with client-specific (P, Q, f) and $(L_{\text{post}}, \rho_{\text{LSI}})$ instantiations. If two or more of these materialize, factoring the cascade out as a template segment becomes attractive.

Why the Fisher-Rao specialization is worth landing even under the small- I restriction. In the AAT-canonical operating regime, Class 3 (Coupled) bias is *moderate* — the goal-conditional reweighting shifts the posterior but does not replace it. Small- I is the honest model of this regime; large- I is the structural-failure limit where the coupled agent’s goal-conditioning overwhelms evidence. Track 2’s universal $C_{FR} = \sqrt{2}$ is the right bound in the operating regime; Track 1 is the fallback under structural conditions where I is not small or where Fisher-Rao is not available.

Connection to logogenic agents. This appendix’s bounds are the quantitative apparatus underlying `#scope-observation-ambiguity-modulation`’s product-form bias bound and `#result-section-ii-survival`’s Class 3 (Coupled) scope-survival analysis. Both segments previously carried the “constant C not computed” caveat; with this appendix landed, C is computed under named conditions and the epistemic tier of the bias bound upgrades to “conditional (exact under H1-H3 or H1+H4).” The logogenic-agents segments retain their client status; the quantitative work moves here.

Findings.

Universal Constant for the Coupled-Agent Bias Bound under Parameterization-Invariance. Brief: When a goal-conditioned agent (like an LLM) processes evidence, its update can be biased by what it’s trying to do, not just what it observes. Prior work could only say “this bias is bounded by some constant times the relevant information leakage” without computing the constant. This result computes it under two specific conditions, and shows that at least one of them — committing to a particular geometric structure on the model space (the (PI) parameterization-invariance axiom) — is necessary in a strong sense: without it, no universal constant exists at all. A transport-inequality track gives $C_{W_2}^2 = 2L_{\text{post}}^2/\rho_{\text{LSI}}$ linear in the goal-conditional information leakage, under log-Sobolev concentration of the observation distribution and Lipschitz-stability of the Bayesian posterior. A Fisher-Rao track gives a universal dimension-free $C_{FR} = \sqrt{2}$ in the small-information regime, conditional on the (PI) axiom. Two structurally wrong derivation routes (Cramér-Rao inversion; rate-distortion inversion) are recorded as honest failure cases.

Impact: Promotes the bias bound from “order-of-magnitude guidance, not a theorem” to a conditional theorem under explicitly named hypotheses, closing the quantitative gap that `#scope-observation-ambiguity-modulation` and `#result-section-ii-survival` previously carried as a caveat. The no-go strengthens the (PI) commitment from useful axiom to load-bearing structural choice and adds a fourth primary instance to the `#disc-additive-coordinate-forcing` meta-pattern. The recorded failures prevent future agents from re-attempting Cramér-Rao or rate-distortion inversions for upper bounds on side-information-induced displacement, and identify transport-inequality machinery (Pinsker, Otto-Villani, Bakry-Émery, posterior-stability) as the correct family. Three forward-looking clients in adjacent appendices

(causal-IB pushforward, misspecification-cost, composition-scope-robustness) are candidate consumers of the same Track 1 cascade, suggesting eventual extraction of a shared #posterior-displacement-template if two or more materialize.

Novelty Claim: *Claim differentiation* on the Lipschitz-posterior + Otto-Villani composition for AAT’s coupled-agent bias bound, plus *claim novelty* on the no-go counterexample showing that universal C in Euclidean-parameter norms cannot exist, which jointly elevates the (PI) axiom from convergent representational choice to load-bearing for theorem-level status. The two upstream tracks each compose textbook information-geometric / transport-inequality machinery (Cover-Thomas chain rule; Otto-Villani 2000; Stuart 2010 Bayesian-posterior stability; Amari-Nagaoka 2000 Fisher-metric expansion); the contribution is the composition for the Class 3 (Coupled) agent bias setting and the no-go that justifies the geometric commitment. (*This finding’s novelty is the supporting quantitative apparatus; the integrated $\kappa \times A$ bias law it supports is in #scope-observation-ambiguity-modulation, where the Pillar-3 verdict applies.*)

Related Work:

Table AB.2

<i>ASF concern</i>	<i>Prior-art language</i>	<i>Relationship / Positioning</i>
KL-to- W_2 transport under log-Sobolev	Otto & Villani 2000 <i>J. Funct. Anal.</i> 173:361–400 (published 2000); Bakry & Émery 1985 curvature-dimension condition	<i>formal antecedent</i> — adopted directly in the Track 1 cascade; standard transport-inequality machinery, cited and used
Lipschitz-stability of Bayesian posteriors	Stuart 2010 <i>Acta Numerica</i> 19:451–559 (published 2010); Hairer, Stuart & Vollmer 2014 <i>SIAM J. Math. Anal.</i> 46(1):415–451 (published 2014)	<i>formal antecedent</i> — adopted directly as (H3); the conjugate-Gaussian case is computed explicitly in §6
Čencov uniqueness of the Fisher metric	Čencov 1982, <i>Statistical Decision Rules and Optimal Inference</i> (published 1982); Ay, Jost, Lè & Schwachhöfer 2017 <i>Information Geometry</i> Theorem 5.1 (published 2017)	<i>formal antecedent</i> — adopted via the (PI) axiom in #scope-agent-identity and used in the Track 2 derivation; the no-go in §4 is the AAT-internal demonstration that this adoption is load-bearing for the bound, not just convenient
KL-to-Fisher second-order expansion	Amari & Nagaoka 2000, <i>Methods of Information Geometry</i> §3.7 Theorem 3.1; Cover & Thomas 2006 §12.5	<i>formal antecedent</i> — supplies the small- I regime sharpness used in Track 2
Information-processing-constraint bounds on decisions	Ortega & Braun 2012 <i>Proc. R. Soc. A</i> (published 2012-04, found 2026-04 via Undermind); Genewein et al. 2015 <i>Front. Robot. AI</i> 2:27 (published 2015-11, found 2026-04)	<i>adjacent literature</i> — addresses a different quantity (decision quality under information constraint rather than belief-update bias under goal-coupling); shares the information-theoretic vocabulary but does not produce a Wasserstein or Fisher-Rao bias bound for coupled-architecture agents
Variational bounds on posterior beliefs in active inference	Friston et al. 2017, “Active inference: a process theory” <i>Neural Computation</i> 29:1–49 (published 2017)	<i>adjacent literature</i> — uses overlapping transport-inequality machinery in a different framework; consistent with the Track 1 derivation but does not require the variational-free-energy commitment (cf. #form-information-bottleneck Discussion)

Novelty assessment. A prior-art search located bounded-rationality and information-cost literatures (Ortega & Braun 2012, Genewein et al. 2015) as the closest decision-theoretic precursors and confirmed no equivalent of the universal- C

derivation under (PI). The verdict (Wholly Novel, Medium confidence) applies to the integrated $\kappa \times A$ law downstream in #scope-observation-ambiguity-modulation; this segment's contribution is narrower (the constant derivation and the no-go) and inherits the same Medium confidence, with a follow-on targeted search recommended on formal posterior-displacement bounds in coupled-architecture agents.

Search Log:

- 2026-04 (*nominally comprehensive*): A comprehensive-search pass (Pillar 3 of the prior-art defense) located bounded-rationality and information-cost literatures (Ortega & Braun 2012, Genewein et al. 2015) as the closest decision-theoretic precursors and confirmed no equivalent of the universal-C derivation under (PI) was surfaced. The verdict (Wholly Novel, Medium confidence) applies to the integrated $\kappa \times A$ law downstream in #scope-observation-ambiguity-modulation; this segment's narrower contribution (the constant derivation and the no-go) inherits the same Medium confidence, with a follow-on targeted search recommended on formal posterior-displacement bounds in coupled-architecture agents.

Working Notes.

- **Causal-discipline cross-link (2026-05-31).** Recognized as the landed $\varepsilon(\kappa)$ bound for the causal-discipline scoping rule of #der-directed-separation (Stage-2 reframe): §1's baseline terminology-aligned from "Class 1 (Separated) scope" to the causally-disciplined $\kappa_{\text{processing}} = 0$ object (held equally by Class 1 by construction and an idealized Class 2 in effect), and the Discussion now names #def-value-object's Q_O independence as the first realized client. No re-derivation — the bound already carried the factor- κ form; this surfaces it as the degradation norm's worked model. The Class 2 (Partial) closed-form refinement ($\kappa_{\text{processing}} \in (0, 1)$) remains open per the Working Note below.
- **Trail.** Full derivation history in spikes/.integrated/spike-bias-bound-constant-C-strengthening-2026-04-24.md. The spike's Track 1 and Track 2 structure maps directly to §§2–3 here; the no-go Attempt E maps to §4; the failed attempts (F1 Cramér-Rao, F2 rate-distortion) map to §5. The spike also flags the shared Pinsker / Otto-Villani machinery with #deriv-variational-sector-condition as a candidate shared-lemma extraction (spikes/spike-kl-to-state-distance-template-extraction-2026-04-24.md Option B); that extraction is contingent on this appendix landing + forward-looking clients materializing.
- **L_{post} for AAT-native cases.** The Lipschitz-posterior stability constant L_{post} is cited from Stuart 2010 / Hairer-Stuart-Vollmer 2014 but not computed for all AAT-relevant cases. The conjugate-Gaussian case of §6 is explicit; general non-parametric Bayesian inverse problems have looser bounds. A future refinement could compute L_{post} for the exponential-family / conjugate-prior cases in 4.3 directly.
- **Large- I Fisher-Rao refinement.** Track 2's $\sqrt{I}$ scaling is the small- I Taylor expansion. Extending to large I requires either compact-manifold assumption or staged composition; flagged for future work.
- **Interaction with the $\mathcal{A}$ factor from #scope-observation-ambiguity-modulation.** The product form $\kappa_{\text{eff}} = \kappa \cdot \mathcal{A}$ interacts with the transport-inequality constants in ways not fully worked out here. If $\mathcal{A}$ is itself an information-theoretic ratio, the $\mathcal{A}$ factor may be absorbable into the I term rather than appearing as a separate multiplicative factor. Flagged for a clarification pass when the logogenic-agents segments are next revised.
- **Non-Bayesian pushforwards.** The Track 1 (T3) posterior-Lipschitz step is stated for Bayesian-posterior pushforwards. Non-Bayesian pushforwards (e.g., post-intervention transformations in the causal-IB candidate) would use the general W_2 -Lipschitz form; the machinery generalizes, though the concrete constants are client-specific.
- **Relation to active-inference bounds.** Active inference (Friston 2017 *Active inference: a process theory*, *Neural Computation* 29) uses KL bounds on posterior beliefs under the free-energy principle; those bounds are formulated in a different framework but share the underlying transport-inequality machinery. AAT's Track 1 derivation is consistent with the active-inference reading but does not require the variational-free-energy-as-master-objective commitment (see 2.2 Discussion for the lineage distinction).

- **Class 2 (Partial) bias-bound analog — open derivation.** This segment derives the bound C for the Class 3 (Coupled) extreme at $\kappa_{\text{processing}} \rightarrow 1$. The corresponding bound for Class 2 (Partial) agents at $\kappa_{\text{processing}} \in (0, 1)$ is open. Both Track 1 (transport-inequality under LSI + Lipschitz-posterior) and Track 2 (Fisher-Rao under (PI) + Čencov + small- I) derive the worst-case envelope; specializing either track to the bounded-coupling regime is the natural starting point. Whether the Class 2 case admits a clean closed-form bound or requires architectural specifics (e.g., per-stage modularity decomposition) is not yet known. Surfaced 2026-05-09 by the GUC rename audit; flagged in `TODO.md` §"Open theory items (MEDIUM)" for cross-referencing.
- **Migration note (2026-05-09 GUC rename):** Class 2 $\leftrightarrow$ Class 3 swap. Pre-2026-05-09: Class 2 = fully merged / fully coupled, Class 3 = partially modular. Post: Class 2 = Partial, Class 3 = Coupled. The segment's canonical claim previously known as the "Class-2 ambiguity bias bound" (fully-coupled case) is now the "Class 3 (Coupled) ambiguity bias bound." Every "Class-2 (fully-coupled)" reference renamed to "Class 3 (Coupled)"; "Class 3 partially-modular" $\rightarrow$ "Class 2 (Partial)"; "Class-1 (modular)" $\rightarrow$ "Class 1 (Separated)." Downstream segments citing the old "Class-2 bias bound" name: see external-citation list in the parent Phase 2 commit report — those segments are scope for later Phase 2/3 work. Removed at candidate stage per `FORMAT.md` Gate 4.
- **Cross-reference to NeurIPS Paper 3.** This segment's two-tracks framework is materially sharpened in NeurIPS 2026 Paper 3 ("How Much Can LLMs Hallucinate? An Upper Bound on Goal-Coupling Displacement", `~/src/neurips/03-llm-hallucinate-bound/`). The paper carries: (i) **Two named tracks with explicit hypothesis sets** at §4 Theorem 4.1 / `#thm-umbrella` — Track 1 (transport-inequality, W_2 metric) recovers the canonical Stuart-school cascade form $C_{T_2} \propto L_{\text{post}}^2 / \rho_{\text{LSI}}$ under (H2'); Track 2 (Fisher-Rao Čencov uniqueness) gives $C = 2$ **globally on the ambient spherical-arc pseudometric** — the operating constant for adversarial / rare-high-KL prompts including jailbreaks and persona injection (via Rényi-1/2 monotonicity + chord-arc on the unit L^2 -sphere; sharp via symmetric N -point witness) — tightening to $C = \sqrt{2}$ **locally** under (R) + (K) + (H4') uniform-locality on goal-conditional slices (first-moment sharp via symmetric two-point witness). (ii) **Chain-rule on post-update law as bridging lemma** at §5 Lemma 5.1 / `#lem-chain-rule`: $\mathbb{E}_G[\text{KL}(P_{M_{\tau+}|e, M_{\tau-}, G} \| P_{M_{\tau+}|e, M_{\tau-}})] = I(G; M_{\tau+} | e_{\tau}, M_{\tau-})$ — the structural identity connecting Bayesian inverse-problems with architectural classification. (iii) **No-go theorem on Euclidean chart norms** at §4 Theorem 4.2 / `#thm-no-go`: chart-rescaling argument $\phi \mapsto a\phi$ scales chart-Euclidean W_2 linearly while leaving $\text{KL}/\text{MI}/d_{FR}/\text{Hellinger}$ chart-invariant; taking $a \rightarrow \infty$ contradicts any candidate fixed $C_0\sqrt{I}$. *Forces (PI) as load-bearing for theorem-level status* and supplies the explicit witness construction for what this segment's Attempt E flagged structurally. The fourth instance of M1 identifiability-floor (named F4 in Paper 3 §4 Lemma 4.2). (iv) **Čencov-uniqueness-and-sharpness** at §4 Theorem 4.3 / `#thm-fr-uniqueness`: under (H1)+(PI)+(R)+(K)+(H4'), $d_{\mathcal{M}}$ is uniquely the ambient Amari-Nagaoka Fisher-Rao spherical-arc distance and $C = \sqrt{2}$ is the unique sharp upper-bound constant. (v) **Class 1 reduction theorem to Stuart-school setup** at §App-D — the Class 1 (Separated) specialization of Track 1's hypothesis space *is* the Stuart-school Lipschitz-posterior cascade (the strict-strengthening from spike A7). The paper's (PI)+(R)+(K) axiom triple — chart-invariance + Riemannian structure + KL second-order matching — is the load-bearing structural commitment for the universal-constant route. See `spikes/neurips-back-integration-2026-05-08.md` §1 Paper 3 entries 1, 2, 3, 4, 5, 7. Standalone-segment proposals under back-integration §3: `#deriv-post-update-chain-rule`, `#deriv-chart-rescaling-no-go`.

Appendix AC

Discussion: Partial-Coupling Pathways — A Sub-Typology Within Class 2

Class 2 (Partial Coupling) per *#der-directed-separation* is named structurally as “some shared infrastructure, some separate pathways” and operationalized as a scalar $\kappa_{\text{processing}} \in (0, 1)$ that aggregates over the un-modeled internal sub-structure. The scalar is distribution-dependent and tracks the *magnitude* of goal-coupling, but not its *kind*: it cannot distinguish the agent that uses goals to *attend* from the one that uses goals to *interpret*, or the agent biased by *what it wants to be true* from the one foreclosed by *what plan it has committed to*, or the additively-biased agent (whose coupling is wrappable by post-hoc debiasing) from the multiplicatively-entangled agent (whose coupling is not). These distinctions are operationally distinct — they recommend different repairs, predict different dynamical signatures, and gate different wrapping regimes.

This segment names the structural sub-space the scalar projects over. Class 2 partial coupling decomposes along three structurally-independent axes — *stage*, *source*, and *form* — yielding a parameterization within which the Class 3 (fully Coupled) corner and the Class 1 (Separated) corner are recovered as the two trivial endpoints. The stage axis tracks which step of the belief-update pipeline takes a G_t argument; the source axis tracks whether the coupling comes from the objective O_t or from the strategy Σ_t ; the form axis tracks whether the coupling acts as separable additive bias on a stage’s output (*content*) or as a non-separable change of the stage’s functional form (*process*). The three axes have distinct operational consequences (Discussion §“Operational consequences” below), and the form axis at Class 2 unifies with the *structure-vs-behavior* refinement at Class 1 — they are the same distinction read at different architectural levels.

The treatment is discussion-grade at the structural-recognition level. The pipeline decomposition is posited as canonical for Bayesian-style updaters and serves as an interpretive lens for monolithic architectures. The substantive structural results that fall out — form-determines-wrappability (a content-form / process-form no-go cut), and the source asymmetry of belief-strategy attractors — are stated qualitatively here and derived in their own segments (*#der-belief-strategy-attractor* for the source-asymmetry result; the form no-go is stated formally in this segment’s Discussion). The combinatorial parameterization is structurally complete but operationally too fine; the load-bearing distinctions live on the (form $\times$ source) projection with stage as a localization tag.

Formal Expression.

Pipeline decomposition of f_M .

Any belief-update function f_M that maintains M_t — whether Bayesian, Kalman, RL with separate world model, or LLM-as-tracker — factors functionally through a small set of operations. The canonical decomposition is

Formulation (partial-coupling-pipeline).

$$f_M = \tau \circ \alpha \circ \lambda \circ \phi \quad (\text{AC.1})$$

with the four stages mapping to AAT machinery already in canon:

Table AC.1

Stage	Operation	AAT correspondent
P1 — Featurization	$\phi : \mathcal{E} \times \mathcal{M} \rightarrow \mathcal{X}$, $x = \phi(e_\tau, M_{\tau-})$	Extracts features from the realized event given current model; in a Kalman tracker, the innovation $\delta = o - \hat{o}$ per #def-mismatch-signal
P2 — Likelihood evaluation	$\lambda : \mathcal{X} \times \mathcal{M} \rightarrow \mathcal{L}$, $\ell = \lambda(x, M_{\tau-})$	Evaluates the likelihood of features under representable hypotheses; in a Kalman tracker, the Gaussian likelihood on δ with innovation covariance S
P3 — Aggregation	$\alpha : \mathcal{L} \times \mathcal{M} \rightarrow \mathcal{M}^+$, $M' = \alpha(\ell, M_{\tau-})$	Combines likelihood with prior to produce new belief; in a Kalman tracker, the gain-weighted update with $\eta^* = U_M / (U_M + U_o)$ per #emp-update-gain
P4 — Consolidation	$\tau : \mathcal{M}^+ \rightarrow \mathcal{M}$, $M_{\tau+} = \tau(M')$	Post-update transformation: storage, normalization, regularization, memory consolidation

Each of (P1)–(P4) can independently take or omit a G_t argument. The Class 1 case has none; the Class 3 case has all; Class 2 is everything between.

The decomposition is posited rather than derived. It is canonical for Bayesian-style belief-updating agents and serves as an interpretive lens for monolithic architectures (where the stages are not architecturally distinct in the forward pass) — the typology’s precision drops the more monolithic the agent.

The (stage $\times$ source $\times$ form) parameterization.

A Class 2 partial-coupling sub-type is a triple $\mathcal{C}_2 = (S, R, F)$ with:

Definition (class-2-sub-type).

- $S \subseteq \{P1, P2, P3, P4\}$, $S \neq \emptyset$ and $S \neq \{P1, P2, P3, P4\}$ — the **stage set** at which goals enter f_M . Empty S recovers Class 1; full S (with all-process form) recovers Class 3.
- $R \subseteq \{O, \Sigma\}$, $R \neq \emptyset$ — the **source set** within $G_t = (O_t, \Sigma_t)$ that acts as the coupling input.
- $F : S \rightarrow \{\text{content}, \text{process}\}$ — the **form** at each coupled stage.

The endpoints recover canon: $S = \emptyset$ is Class 1; $(S, R, F) = (\{P1, P2, P3, P4\}, \{O, \Sigma\}, \text{process})$ is Class 3 (fully Coupled). The scalar $\kappa_{\text{processing}}$ is the projection of $\mathcal{C}_2$ onto the magnitude axis.

Form classification — content vs process.

A stage operation $\xi(\cdot; G)$ is **content-form coupled** if there exist functions ξ^0 (goal-blind) and b_ξ (bias) such that

Definition (coupling-form).

$$\xi(u; G) = \xi^0(u) + b_\xi(G; u) \quad (\text{AC.2})$$

with b_ξ identifiable from a probing protocol that varies G at fixed u (specifically: a reference goal G_0 exists with $b_\xi(G_0; u) = 0$, or b_ξ is identifiable up to a u -dependent constant by varying G across finitely many probes).

A stage operation is **process-form coupled** if no such decomposition exists — the operation's *functional form* depends on G , not just its output. The multiplicative case $\xi(u; G) = \xi^0(u) \cdot h(G; u)$ is the most common process form (multiplicative attention coupling, gain modulation, precision weighting); the compositional case $\xi(u; G) = \xi^\dagger(u, G)$ with no separable factor is the starkest form.

The classification is *probing-protocol-dependent*: an analyst with reference-goal access may classify as content where an analyst without that access classifies as process. This is the same epistemic situation as the existing $\hat{\kappa}_{\text{processing}}$ behavioral estimator — the form classification inherits the same probe-dependence.

Form-determines-wrappability — the central operational claim.

For a Class 2 sub-type with coupling triple (S, R, F) :

Derived (content-process-wrappability, conditional on probing model).

Positive half (content-form is post-hoc-wrappable). Stages with content-form coupling admit an *external debiasing wrapper* that requires only behavioral-probing access (no internal pipeline access). The wrapper estimates $b_\xi(G; u)$ from probes that vary G at fixed u , then subtracts the estimated bias in normal operation. The wrapped output converges to $\xi^0(u)$ as the estimator becomes consistent. This corresponds to the W_2 (partial-wrapping) regime at the agent level — “the query boundary still passes G_W to the component” with typed-output structuring extracting a goal-blind belief-update.

Negative half (process-form is not post-hoc-wrappable). Stages with process-form coupling do not admit such an external wrapper. The argument is identifiability-theoretic: under process-form there is no G_0 making $h \equiv 1$ (or, more generally, no u -independent reference goal recovering ξ^0); behavioral probes that vary G reveal goal-correlated differences but cannot collapse them to a G -invariant honest output. Repair requires either *stage replacement* (substituting a goal-blind version of the stage in the pipeline — requires substitutability access at the stage's input/output boundary) or *full-agent wrapping* (the W_1 strict-wrapping regime, operating at the whole-agent boundary).

The form distinction is therefore *load-bearing*, not ornamental — it gates the wrapping regimes available for coercion to Class 1.

Stage-cascade propagation.

Goal-coupling at stage P_k contaminates all downstream stages $P_{k+1}, \dots, P_4$ even when those downstream stages are individually goal-blind, because their inputs are functions of G through the cascade. The *naming* of a sub-type by its primary coupling stage refers to where the coupling originates; the *effect* propagates forward through the pipeline. Minimal repair (per stage-localization-of-repair, below) is therefore at the *most upstream coupled stage*, not at the most downstream symptom.

Derived (cascade-propagation).

Stage-localization of repair.

If pipeline access is available at stages $P_k \in S$ — the wrapper can intercept the stage’s input and supply its output to P_{k+1} — then substituting goal-blind operations P_k^0 at every coupled stage produces a wrapped agent that is Class 1 at the wrapper level. The construction has a capability cost (the substitute operations may be weaker than the agent’s native ones) and a Brooks’s-Law tempo overhead (more pipeline calls per macro-step), but the structural class-coercion is exact.

Derived (stage-localization-of-repair, conditional on pipeline access).

For monolithic architectures without pipeline access, only full-agent wrapping (the construction of `#der-class-coercion-via-wrapping`) is available; stage-level repair does not apply.

Composition with the leakage locus. Goal contamination’s *effect on belief* is confined to the Fisher null space $\ker J_\tau$ of the observation given the current latent state (the leakage-locus result, when promoted to its own segment, will be the formal statement). This locus is *universal across Class 2 sub-types* — it does not depend on which pipeline stage the coupling enters. The sub-type determines the *functional form* of the displacement within the locus:

- Content-form sub-types produce displacement linear in G (the $\Delta\mu = \Lambda_0^{-1}g(G)$ shape in the linear-Gaussian instantiation, with $g(G)$ a linear function of the goal-tilt parameter).
- Process-form sub-types produce displacement with a G -dependent gain matrix or a G -dependent covariance structure — non-linear in G .
- Σ -source sub-types compound the per-step displacement across the closed-loop dynamics; under multiplicative process form, self-stabilizing belief-strategy attractors arise (`#der-belief-strategy-attractor`).
- O -source sub-types produce exogenously-bounded tilt magnitude (because O_t revises only when forced, per `#der-orient-cascade`).

Wrapping-regime correspondence. The $W_0 / W_1 / W_2$ wrapping regimes of `#der-class-coercion-via-wrapping` are the agent-level analogs of the stage-level form distinction:

Table AC.2

<i>Wrapping regime</i>	<i>Stage-level analog</i>	<i>Class achieved at wrapper</i>
W_0 (no wrapping)	n/a — the agent stands as-is	Whatever the un-wrapped agent is
W_2 (partial wrapping)	Content-form at the agent level — wrapper passes G to component, post-hoc structures the response	Class-1-by-behavior (per <code>#der-directed-separation</code>)
W_1 (strict wrapping)	Process-form-with-pipeline-access — wrapper structurally substitutes goal-blind execution	Class-1-by-structure

This unifies the *structure-vs-behavior* distinction at Class 1 with the *content-vs-process* form distinction at Class 2: they are the same axis read at different architectural levels. The structure / behavior distinction at the Class 1 cell is the agent-level shadow of a sub-classification axis that runs all through Class 2.

Sub-type matching refines the wrapping-regime recommendation: a Class 2 agent with content-form coupling only is sufficient with W_2 ; a Class 2 agent with process-form coupling at any stage requires W_1 or remains effectively Class 2 under W_2 ; a Class 2 agent with Σ -source coupling may need a Σ -channel-suppressed W_1 even with pipeline access, because strategy-context can leak through stateful component channels independent of the query. The call-boundary condition

under which W_1 's structural guarantee survives such stateful leakage is the (C2') condition of #disc-w1-structural-bound-boundary: when goal-correlated state crosses the call boundary, W_1 's structural bound is unavailable and only a behavioral bound remains.

Epistemic Status.

Discussion-grade at the structural-recognition level. The pipeline decomposition is a *formulation choice* — canonical for Bayesian-style updaters and well-motivated by AAT's existing machinery, but not forced by AAT canon. The (stage $\times$ source $\times$ form) parameterization is derived from the choice. The form distinction is *robust qualitative* at the operational level (the content / process distinction has a clean identifiability argument); the formal definition is probing-protocol-dependent.

The two substantive structural consequences carry their own tiers:

- **Form-determines-wrappability** (this segment, Discussion §): *robust qualitative* — the identifiability argument is exact under the formal non-separability definition; the operational dependence on probing-protocol availability is honest scope.
- **Source asymmetry / belief-strategy attractors** (#der-belief-strategy-attractor): *conditional* — exact under the linearized dynamics and the posited multiplicative form for Σ -source aggregation-stage coupling; robust qualitative beyond.
- **Composition with leakage locus**: *derived* under the linear-Gaussian instantiation; robust qualitative more generally (it composes existing structural results).
- **Wrapping-regime correspondence**: *robust qualitative* — the structural correspondence is exact, but boundary cases (partial-information wrappers between W_1 and W_2) require per-case analysis.

The full $42 \times 2^{|S|}$ -cell parameterization is structural completeness; operationally the load-bearing distinctions live on the coarsening to *upstream-most coupled stage $\times$ source-set $\times$ form*, which gives roughly 30 cells, of which empirically about a dozen are populated by recognizable phenomena (canonical phenomena placed in Discussion §"Canonical phenomena" below).

Discussion.

Why the form axis is the operationally-load-bearing one. Of the three axes — stage, source, form — the form axis is the one that gates wrappability: it determines whether external post-hoc debiasing suffices or whether stage replacement / full-agent wrapping is required. The source axis is the one that gates *dynamics*: Σ -source coupling closes a feedback loop through the orient cascade and can produce self-stabilizing attractors that O -source coupling cannot. The stage axis is the one that gates *architectural location of repair*: it determines where in the pipeline the substitution or debiasing must be applied. These three operational outputs from the typology — wrappability, attractor possibility, repair location — each correlate with one axis, which is why the three-axis decomposition is more informative than the scalar.

Unification with Class 1's structure-vs-behavior refinement. The Class 1 cell of #der-directed-separation admits a refinement: *Class-1-by-structure* (no G in the belief-update query — structural commitment by type signature) vs *Class-1-by-behavior* (the query passes G but the component is instructed not to use it — separation by compliance). The form distinction at Class 2 is the same axis: *content-form* is the behavior-side (the operation accepts G but its goal-shaped component is post-hoc-subtractable); *process-form* is the structure-side (the operation's functional form depends on G inseparably, requiring substitution). Read together, the structure / content side and the behavior / process side describe two distinct *modes* of goal-coupling — one in which the coupling can be neutralized by *not asking* (or by subtracting after asking), and one in which the coupling is constitutive of the operation.

Why the Σ -source / O -source asymmetry matters. The orient cascade (#der-orient-cascade) has Σ updated as a function of (M, O) and O revised only when forced. This makes Σ *endogenous* to M via the cascade — Σ updates from

the new M that the agent's f_M produces — but leaves O *exogenous* to M in steady state. The asymmetry feeds back: a Class 2 agent with Σ -source coupling at any stage of f_M has the structure $M \rightarrow \Sigma \rightarrow f_M \rightarrow M$ closing a loop, while the same coupling sourced from O has no such closure (because O does not update from M per step). The result is that Σ -source coupling under multiplicative process-form admits *self-stabilizing fixed points* in which M stays misaligned with the environment indefinitely — the formal sunk-cost attractor of #der-belief-strategy-attractor. O -source coupling produces *bias* but not *runaway commitment*.

Canonical phenomena placed in the typology. The cells the typology produces are operationally meaningful:

Table AC.3

<i>Phenomenon</i>	<i>Stage(s)</i>	<i>Source</i>	<i>Form</i>	<i>Repair regime</i>
Goal-directed attention	(selection — pre- f_M)	O or Σ	—	None — formally allowed by #der-directed-separation scope condition
Motivated reasoning (identity-driven)	P1 + P2	O	Content (moderate) $\rightarrow$ Process (strong)	W_2 moderate; W_1 + identity-suppression strong
Sunk-cost commitment	P3	Σ	Process (multiplicative)	W_1 + Σ -suppression; admits attractors per #der-belief-strategy-attractor
Identity-protective consolidation	P4	O	Process (compositional)	W_1 + storage-protocol externalization
Frame coupling	P1	$O + \Sigma$	Process (compositional)	Multi-frame composition (not coercion to Class 1)
Wishful thinking	P3	O	Content (mild) $\rightarrow$ Process (strong)	W_2 mild; W_1 strong
Affect / urgency cascade-bypass	(cascade-bypass)	n/a (not stage-coupling)	n/a	External pacing scaffolding (per #disc-adversarial-coupling-pressure affect mechanism)
Attention-mediated transformer-LLM	P1+P2+P3	$O + \Sigma$ (prompt-mixed)	Process (compositional)	Full Class-3 — W_1 wrapping construction, no stage repair without internal access

The transformer-LLM case lands at the corner ($S = \{P1, P2, P3\}$ or $\{P1, P2, P3, P4\}$, $R = \{O, \Sigma\}$, $F \equiv$ process), which is the parameterization's structural corner near Class 3 (Coupled). The typology recovers canon at its endpoints. Classical confirmation bias is a *separate* failure mode: it is structurally $M_t^{\text{prior}} \rightarrow f_M$ self-coupling, not $G_t \rightarrow f_M$ coupling, and is allowed by #der-directed-separation's scope condition (which has $M_{\tau-}$ as an argument); it lives outside the (S, R, F) parameterization, except in the *trajectory-coupling* case where the prior has been historically shaped by G -coupled updates and the cumulative effect carries effective G -content. Trajectory coupling is a separate per-trajectory axis from the per-step typology.

Connection to the three operations on modularity state. The M4 meta-segment #disc-modularity-state-dynamics names three operations — truthification, strategic self-coupling, adversarial coupling pressure — that change modularity state. The typology here is *orthogonal* to those operations: the operations describe *transitions* through the parameterization (truthification reduces S or simplifies F ; strategic self-coupling expands S and shifts F toward process; adversarial pressure expands S from the outside, targeting specific cells — identity-binding targets P1+P2 / O ; sunk-cost-engineering targets P3 / Σ); the typology describes the *static structural sub-space* on which the operations act. Each adversarial-coupling-pressure mechanism corresponds to a specific cell, and the three-mechanism table is best read as the projection of the typology's static structure onto the *externally-driven-decreasing* leg of M4's three-operation pattern. M4 names the operations; this segment names the structural state-space.

Operational consequences — refined behavioral estimator. The existing $\hat{\kappa}_{\text{processing}}$ behavioral estimator (per #der-directed-separation §”Empirical estimator”) aggregates over the typology. A refined estimator would distinguish:

- *Stage*: probing intermediate representations (where mechanistic-interpretability access exists) localizes the upstream-most coupled stage.
- *Source*: probing under O varied at fixed Σ vs Σ varied at fixed O separates the source set.
- *Form*: probing distinguishes mean shifts (content signature) from covariance shifts or functional-form changes (process signature) in the agent’s responses across goal-variants.

The refinement targets *which sub-type* the agent inhabits, not just the aggregate magnitude. Per the composition with the leakage locus, all such probing should be confined to $\ker \mathcal{J}_\tau$ — the directions the observation does not identify — which is where the displacement lives.

Honest scope. The typology is most precise for *modular* Class 2 architectures (biological cortex per the canonical example; hybrid AI systems with separable preprocessing / inference / aggregation modules). For monolithic architectures (transformer LLMs, end-to-end-trained policies), the stage decomposition is *interpretive* — the stages are not architecturally distinct in the forward pass; the typology serves as a lens for mechanistic-interpretability investigations rather than as a structural classification. The form classification is *probing-protocol-dependent*; analysts should specify which probes were available. Content-form wrapping gives *Class-1-by-behavior*, not honesty: under gauge freedom in the identifiability, the wrapper produces a G -invariant version of the agent’s belief, up to an unknowable additive constant.

Findings.

Partial-Coupling Sub-Typology (Stage $\times$ Source $\times$ Form). **Brief:** The Class 2 (Partial Coupling) label in AAT names “some shared infrastructure, some separate pathways” between belief-update and goal-state processing, but the existing scalar measure $\kappa_{\text{processing}}$ collapses a rich internal sub-structure into a single magnitude number. Look more carefully and three structurally-independent axes emerge: *which stage* of the belief-update pipeline takes a goal argument (attention / interpretation / aggregation / storage); *which source* of goal-state acts (the agent’s objective, or its committed strategy); and *which form* the dependency takes (additive bias subtractable by post-hoc debiasing, vs functional-form change requiring internal substitution). Each axis carries a distinct operational consequence: the form axis gates whether external wrapping can coerce the agent back to Class 1; the source axis gates whether self-stabilizing belief-strategy attractors are possible; the stage axis gates where the repair must be applied. The Class 3 (fully Coupled) limit is the corner where all stages couple, both sources act, and the form is process-throughout; Class 1 is the opposite corner; Class 2 is the structured space between.

Impact: Refines the operational meaning of the Class 2 label from “uncomfortable middle category” to “structured sub-space with specific repair regimes per cell.” Unifies the *Class-1-by-structure* vs *Class-1-by-behavior* refinement (currently sitting at the Class 1 cell of #der-directed-separation) with the form-axis of the Class 2 sub-typology — they are the same distinction read at different architectural levels. Sharpens wrapping-regime selection (W_1 vs W_2) by sub-type rather than by Class label alone. Maps the three adversarial mechanisms named in #disc-adversarial-coupling-pressure (identity-binding / affect-urgency / sunk-cost-engineering) onto specific cells of the typology, with the affect/urgency case recognized as outside the static (stage $\times$ source $\times$ form) parameterization (it is a cascade-bypass, a separate failure mode). Provides the static-structural complement to the M4 meta-segment #disc-modularity-state-dynamics’s three operations on modularity state — the operations describe transitions through the parameterization the typology names.

Novelty Claim: *Claim recognition* of the (stage $\times$ source $\times$ form) sub-typology as the structural complement to Class 1’s structure-vs-behavior refinement. *Claim differentiation* on the wrapping-regime correspondence: which wrapping regime suffices for Class 2 $\rightarrow$ Class 1 coercion is determined by the sub-type’s form, not just by the Class label. The integration with the leakage-locus result (locus universal in $\ker \mathcal{J}_\tau$; typology determines functional form within the

locus) and the source-asymmetry implication (#der-belief-strategy-attractor) are the load-bearing structural consequences.

Related Work:

Table AC.4

<i>ASF concern</i>	<i>Prior-art language</i>	<i>Relationship / Positioning</i>
Active inference / FEP fused goal-perception	Friston 2010+ corpus; covered in <code>ref/Prior_art_for_AAT_directed_separation.md</code> Pillar	<i>adjacent literature</i> — the unified variational machinery is structurally process-form at all stages (the Class 3 case); the <i>partial-separation sub-typology</i> is not addressed in this literature
Motivated reasoning / cultural cognition	Kunda 1990 <i>Psych Bull</i> ; Kahan cultural-cognition project	<i>not yet searched</i> — primary novelty target for the source asymmetry and form-classification of motivated cognition
Identity economics & motivated beliefs	Bénabou-Tirole 2002, 2011 <i>QJE</i> ; Akerlof-Kranton 2000 <i>QJE</i> , 2010	<i>not yet searched</i> — primary target for whether the O -source vs Σ -source asymmetry has an economics-side antecedent
Dual-process theory	Evans 2008; Kahneman 2011; Stanovich 2009	<i>not yet searched</i> — coarse Class 1 / Class 3 distinction is informally present; the sub-typology of partial coupling is not the literature's frame
Self-deception philosophical taxonomy	Mele 2001 <i>Self-Deception Unmasked</i> ; von Hippel & Trivers 2011 <i>BBS</i>	<i>not yet searched</i> — possible direct anticipation on the philosophical sub-classification side
Goal-shielding	Shah, Friedman & Kruglanski 2002 <i>JPSP</i>	<i>not yet searched</i> — candidate empirical instantiation of $P3 / \Sigma$ / process-form coupling

Search Log:

- 2026-05-22 (*intuition-only*): The active-inference / Pearl-blanket / Friston-blanket prior-art landscape is already covered in `ref/Prior_art_for_AAT_directed_separation.md` (Pillar-3 search) and addresses the unified variational machinery as the Class 3 (fully Coupled) case; the *partial-separation sub-typology* this segment develops is not in that prior-art landscape. The cognitive-psychology / motivated-reasoning / identity-economics / dual-process / self-deception literatures are *unsearched*; targeted Pillar search recommended before any tier upgrade. The form distinction (content vs process) likely has analogs scattered across literatures (most explicitly in Bénabou-Tirole's identity-vs-instrumental belief framing and in dual-process system 1 / system 2 separation); the source-asymmetry result is plausibly partially in the sunk-cost / commitment-cascade empirical literature implicitly but not as a derived structural theorem. The combined (stage $\times$ source $\times$ form) parameterization in AAT-internal vocabulary is plausibly novel under the search-depth conducted.

Working Notes.

- **Promotion path requires two sub-spikes.** (a) Pillar-style targeted prior-art search across motivated-reasoning, identity-economics, dual-process, and self-deception literatures; (b) derivation of the multiplicative form for Σ -source aggregation-stage gain $K(\Sigma)$ from utility-cost analysis of belief-revision-under-strategic-commitment (this gates #der-belief-strategy-attractor's tier from *conditional* to *exact*). Both are in `TODO.md` per the queue maintenance.
- **Composite-level sub-type inheritance.** The composite-level class-inheritance table in `#der-directed-separation` (axis-decomposed by routing $\times$ substrate $\times$ dynamic-regime) tracks Class (1/2/3) inheritance. Extending the table to track Class 2 *sub-type* inheritance under composition is a substantive follow-on: given sub-agents with sub-types (S_i, R_i, F_i), what sub-type does

the composite occupy under routing R and substrate σ ? Not on the critical path for this segment; recorded as a candidate sub-spike for the composition machinery to extend.

- **Trajectory coupling as separate axis.** The per-step (stage $\times$ source $\times$ form) parameterization does not cover *trajectory* coupling — cumulative goal-content accumulating in M_t 's prior across history through repeated small Class 2 couplings, then amplified by goal-blind processing of new evidence. This is a real failure mode (and arguably the operational route to confirmation-bias cascades) but lives in a per-trajectory dynamics layer rather than the per-step structural typology. A separate spike treating trajectory coupling would parallel this segment's per-step analysis; not pursued here.
 - **Refined behavioral estimator — construction.** §Discussion sketches a refined estimator that probes along $\ker \mathcal{J}_\tau$ and distinguishes mean (content) from covariance/functional-form (process) signatures. Actual estimator construction is a separate work item; the segment names the targets but does not construct the estimator.
 - **Brief is not yet at Feynman criterion.** The current Brief reaches for the structural-recognition framing rather than an everyday physical analog. A candidate analog under consideration: the difference between (a) an engineer reading instruments and writing a goal-blind report, then a separate engineer reading the report and deciding what to do, versus (b) a single engineer reading instruments with the decision already in mind, where what they notice depends on what they're trying to accomplish — the first is Class 1, the second Class 3, and Class 2 is the messy middle where the typology says *which* part of the data-reading process is being shaped by the goal, *whether* the shaping is what-the-engineer-wants or what-plan-they're-committed-to, and *whether* the shaping is a measurable bias (subtractable) or a wholesale reframing (not subtractable). The bathtub-equivalent has not been written; the engineer-and-report analog is a working candidate but its isomorphism to the load-bearing structure (especially the source-asymmetry attractor and the form-determines-wrappability cut) needs verification.
 - **Connection back to #def-agent-spectrum Working Note.** The Working Note in #def-agent-spectrum ("The one genuinely-new thread: where the dangerous coupling concentrates") points at the leakage locus as a *where-in-the-inference* localization. This segment is the *through-what-stage-of-the-inference* complement. Together with the leakage-locus result they answer the two complementary questions: *where in state-space the effect lives* (leakage locus) and *what functional form it takes there* (this typology). When the leakage-locus result lands as its own appendix segment, the cross-reference should be made explicit on both sides.
-

Appendix AD

Discussion: The W_1 Structural-Bound Boundary

The W_1 structural leakage bound of `#der-class-coercion-via-wrapping` exists if and only if the wrapped component carries no goal-correlated state across the call boundary (condition (C2')); otherwise the goal-leak routes through a channel the wrapper cannot observe or condition on, and only a behavioral (W_2 -type) bound remains. The discontinuity at the boundary is a *certifiability* discontinuity, not a behavioral one: the leakage itself is continuous — and quadratically flat — in the degree of (C2')-violation, while what snaps at the boundary is the availability of the structural certificate.

Formal Expression.

Setup: the two channels and the conditioning that distinguishes them. Recall from `#der-class-coercion-via-wrapping` that the W_1 regime issues a structurally goal-blind query q_M to a black-box component A , then updates the wrapper belief M_W from the response $A(q_M)$. Two distinct goal variables are in play: the wrapper's internal register G_W (which q_M structurally does not take as an argument; Theorem 1 closes this *processing* channel exactly), and the latent operator goal G^{op} inferable from query content (the *selection* channel bounded by $\kappa_{W_1}^{\text{sel}} = I(A(q_M); G^{\text{op}}) \leq I(q_M; G^{\text{op}})$). The structural bound is a property of the wrapper's query-selection policy because, under (C2'), $G^{\text{op}} \rightarrow q_M \rightarrow A(q_M)$ is a Markov chain — the goal can reach the response *only* through the query the wrapper chose.

The no-go: statefulness reopens an unobservable channel.

Drop (C2'). Let A carry hidden state S — a conversation KV-cache, a retrieval buffer, adaptive context — that is correlated with G^{op} through shared history (e.g. a prior goal-conditioned call in the same macro-step touched S). Then the response can depend on the goal *through* S , and this dependence survives conditioning on the query:

Discussion (w1-structural-bound-no-go).

$$I(A(q_M); G^{\text{op}} | q_M) \geq 0, \quad \text{and generically } > 0 \text{ when } S \not\perp G^{\text{op}}. \quad (\text{AD.1})$$

Conditioning on q_M no longer closes the path, because the leak does not flow through q_M — it flows through S . The conditioning that *would* close it, $I(A(q_M); G^{\text{op}} | q_M, S) = 0$, is unavailable to any wrapper-level analysis: S is internal to the black-box oracle and the wrapper cannot observe or condition on it. The Markov chain $G^{\text{op}} \rightarrow q_M \rightarrow A(q_M)$ is broken — there is now a second edge $G^{\text{op}} \rightarrow S \rightarrow A(q_M)$ that bypasses the query — so the data-processing argument that yields $\kappa_{W_1}^{\text{sel}} \leq I(q_M; G^{\text{op}})$ no longer holds.

No-go statement. For a component carrying goal-correlated cross-call state, *no bound on the belief-channel goal-leakage is derivable from the wrapper's query-selection policy alone*. The structural W_1 bound is available if and only if (C2') holds. When (C2') fails, the leak flows through an unobservable, unconditionable channel, and only a *behavioral* bound — estimated by monitoring $I(A(q_M); G^{op} \mid q_M)$ in the responses — remains, exactly as in the W_2 regime and for the same structural reason (an unconditionable goal-correlated channel). This locates the structural-versus-behavioral split not only in *wrapper design* (where to place the separation commitment — the $W_0/W_1/W_2$ regime choice) but *also* in a *component property*: statefulness across the call boundary forces a behavioral-only bound regardless of how the wrapper is built. The condition is sharp — it is not online weight-adaptation that breaks the structural bound (a frozen-weights LLM adapts nothing) but mere goal-correlated state persistence across the call boundary, which a frozen-weights LLM with a conversation cache exhibits.

The discontinuity is in the certificate, not in the behavior. The boundary at (C2') is best read as a *certifiability* discontinuity. Parameterize the (C2')-violation by ε , the goal-correlation retained in the cross-call state ($\varepsilon = 0$ is exact statelessness-about-the-goal; $\varepsilon > 0$ is a sliver of goal-correlated memory). Two facts hold at the boundary, and they point in opposite directions:

- **The behavior is continuous.** The actual goal-information reaching the belief is a continuous function of ε through $\varepsilon = 0$ — no jump, no threshold. Two wrappers differing only in ε leak amounts that differ by an amount shrinking with ε ; they cannot be told apart by watching their belief updates.
- **The certificate's validity is a step.** At $\varepsilon = 0$ the design-side certificate — *the goal is not in the query, the component is stateless about the goal, therefore the belief update cannot see the goal* — is true. At any $\varepsilon > 0$ that same certificate is still computable and still asserts the same thing, but is now *false*: the real leak is positive and took the unobservable path. The certificate's *value* (a flat selection-channel reading of 0 under goal-blind selection) is continuous across the boundary; its *truth* is the step function. Once $\varepsilon > 0$ the only honest instrument is behavioral monitoring.

So the wall is in the epistemics, not in the world: the leakage the structural certificate would have bounded is continuous in the (C2')-violation, and the discontinuity is in what is *provable*, not in what the agent *does*. This pre-empts the natural objection — *surely a tiny bit of state can only leak a tiny bit?* — by agreeing with it and relocating the bite: the leak is indeed tiny, but the structural guarantee is gone entirely.

What is solid versus suggestive. Two parts of the picture are exact and carry the no-go:

1. **The behavioral identity at the limit.** At exact statelessness-about-the-goal the structural certificate's zero reading is the truth (the leak is zero); the Markov chain holds and $\kappa_{W_1}^{sel} \leq I(q_M; G^{op})$ is valid.
2. **The certificate-validity step.** The structural certificate is available iff (C2') holds; at any (C2')-violation it asserts a false zero and only the behavioral monitor remains valid. This is forced by the unobservability of S and does not depend on the magnitude of the violation.

What is *suggestive but not proven in general* is the global *continuity-and-flatness* of the leakage past the boundary. In an exact-enumeration toy (single bit, single call, a state-bends-the-answer mechanism) the leak is monotone, smooth, and $\Theta(\varepsilon^2)$ — flat to first order — throughout the violation range. That is moderately strong evidence for the continuity reading and against a behavior-discontinuity reading, but it is a deliberately simple toy: a component whose retained state *amplifies* (state that controls the response gain, or multi-step accumulation across many calls) could in principle exhibit super-linear or threshold leakage, and that is not ruled out. The continuity claim is therefore stated as *robust qualitative* — the certifiability step is the exact load-bearing content; the smoothness of the leak past the boundary is the toy-suggested, not-fully-general, companion.

Epistemic Status.

Robust qualitative. The no-go itself — *the structural W_1 bound is available iff (C2') holds; under goal-correlated cross-call state the belief-channel leak is unconditionable at the wrapper level and only a behavioral bound remains* — is exact: it follows directly from the unobservability of the component's internal state S and the consequent breaking of the $G^{\text{op}} \rightarrow q_M \rightarrow A(q_M)$ Markov chain. The certifiability-discontinuity framing (behavior continuous, certificate-validity a step) rests on (a) the exact behavioral identity at the (C2') limit and (b) the certificate-validity step, both solid. The *general* continuity-and-flatness of the leakage in the degree of (C2')-violation is *suggestive, not proven*: it is exhibited cleanly in an exact-enumeration toy ($\Theta(\varepsilon^2)$ flatness), but a component with an amplifying state mechanism could in principle break smoothness, and that case is not closed. Hence the segment as a whole sits at robust qualitative rather than exact: the boundary's *existence and location* are exact; the boundary's *shape past it* is qualitatively characterized.

Max attainable: robust qualitative for the continuity-shape claim absent a general continuity proof; the no-go's existence-and-location is already exact.

Discussion.

Why this is a sharpening, not a softening, of W_1 . The W_1 regime's structural guarantee is not undercut by this boundary — it is given a precise domain of validity. Theorem 1 of `#der-class-coercion-via-wrapping` (exact directed separation under (C1)–(C3)) is untouched; the selection-channel bound is untouched within (C2'). What this segment adds is the honest statement of *when the structural bound stops existing* and *what replaces it* (a behavioral bound), together with the recognition that the replacement is forced by a component property orthogonal to wrapper design. A no-go with an explicit boundary characterization is more useful than a structural bound asserted without a stated domain.

Practical reading. The $\Theta(\varepsilon^2)$ flatness of the toy carries a counter-intuitive design implication: a near-stateless component is *operationally* almost as safe as an exactly-stateless one (a little cross-call memory leaks very little goal-information), but it is *certificationally* not safe at all (the structural guarantee is lost entirely the instant statelessness fails). The value of enforcing exact (C2')-statelessness across the call boundary is concentrated almost entirely in what it lets you *prove*, not in the marginal leakage it prevents. When deciding whether the engineering cost of true call-boundary state-resetting is worth it over “mostly reset,” the thing being bought is a proof, not a meaningful behavioral delta.

Relation to the Class 2 sub-typology. This boundary is the agent-level companion to the Σ -channel observation in `#disc-partial-coupling-pathways`: a Σ -source coupling in a stateful component can undermine W_1 's structural commitment via strategy-context that persists across calls, which is precisely a (C2')-violation. The sub-typology names the component-internal form of the coupling; this segment names the call-boundary condition under which W_1 's structural certificate survives it.

Relation to the action-side bounded-signaling assumption. The unobservable-channel structure here is the belief-side twin of the action-side bounded-signaling assumption named in `#der-directed-separation` — both are “a goal-correlated channel the formalism does not see.” Whether they are the same phenomenon viewed from two sides is left open.

Findings.

The W_1 Structural-Bound Boundary (Certifiability Discontinuity at (C2')). **Brief:** Imagine a scaffold that keeps a language model honest about its beliefs by only ever asking it goal-blind questions — “summarize this diff,” never “summarize this diff so we can ship the fix.” The promise of strict wrapping (W_1) is that you can *certify from the design alone* that the operator's goal cannot bend the belief update: the goal is structurally not an input to the belief-side query. That promise holds exactly when the component forgets everything about the goal between the goal-blind question and the goal-conditioned one. The moment the component keeps a sliver of goal-touched memory across that boundary — a conversation cache, a retrieval buffer — the goal can sneak into the answer through the memory, on a path the scaffold cannot see, and the design-side certificate becomes false even though it still computes the same reassuring number. The

striking part is *how* it fails: the actual amount of goal that leaks is continuous and, near the boundary, quadratically tiny — you can get a near-stateless component arbitrarily close in behavior to a stateless one. What snaps, all-or-nothing, is not the behavior but the *provability*: you lose the entire structural guarantee for an almost-invisible behavioral cost. You were buying a proof, not a meaningful safety margin.

Impact: Locates the structural-versus-behavioral leakage split of `#der-class-coercion-via-wrapping` in a *component property* (goal-correlated cross-call statefulness) independent of wrapper design, complementing the $W_0/W_1/W_2$ regime choice that locates it in wrapper construction. Promotes the Σ -channel-suppressed- W_1 Working Note of the wrapping segment to a named, load-bearing condition (C2') and gives the precise reason it is needed. Frames the no-go as a certifiability discontinuity — sharpening the CS-norm practice of characterizing *what kind* of boundary a no-go draws — and supplies the practical corollary that exact call-boundary statelessness buys a certificate, not a behavioral delta.

Novelty Claim: *Claim recognition* that the W_1 structural leakage bound's availability is governed by a sharp component-side condition (no goal-correlated cross-call state), with the boundary characterized as a discontinuity in the *validity of the structural certificate* rather than in the agent's leakage behavior — the leakage being continuous, and second-order flat, in the degree of condition-violation.

Related Work: - Cover & Thomas, *Elements of Information Theory* (2nd ed., 2006), data-processing inequality (published 2006, found 2026-05-31) — *formal antecedent* — the DPI along $G^{\text{op}} \rightarrow q_M \rightarrow A(q_M)$ is what holds under (C2') and breaks when an unobservable S adds a bypassing edge. - Pearl, *Causality* (2nd ed., 2009), d-separation / unobservable-confounder structure (published 2009, found 2026-05-31) — *adjacent literature* — the broken-Markov-chain reading (a hidden S confounding G^{op} and $A(q_M)$ that the wrapper cannot condition on) is the causal-graph idiom for the no-go.

Search Log: - 2026-05-31 (*intuition-only*): No dedicated literature search for the certifiability-discontinuity framing as such; the underlying tools (DPI, unobservable-confounder structure) are standard and cited. A targeted search for prior art on “certifiable goal-blindness under stateful components” / structural-vs-behavioral verifiability boundaries in scaffolded-agent design has not been conducted.

Working Notes.

- Reasoning trail (filed): the correction that produced this no-go and the certifiability-discontinuity framing are in `spikes/.integrated/spike-w1-leakage-vacuity-2026-05-31.md` and `spikes/.integrated/spike-w1-w2-boundary-intuition-2026-05-31.md` (cycle narrative: CHANGELOG 2026-05-31).
 - The exact-enumeration sim exhibiting the $\Theta(\epsilon^2)$ flatness (`spikes/sim-w1-w2-boundary.py`, figure `spikes/w1-w2-boundary.png`) sits unintegrated in `spikes/`, blocking on the sim-integration process (G3 in `INTEGRATION-CLEANUP-TODO.md`).
 - **Open: amplifying-state mechanisms.** The general continuity-and-flatness of the leak past the boundary is toy-suggested, not proven. A component whose retained state controls the response gain, or accumulates across many calls, could exhibit super-linear or threshold leakage. Closing this would lift the continuity-shape claim from robust qualitative toward exact; it connects to the M4 “amplification under high component fidelity” thread in `#disc-modularity-state-dynamics` and the accuracy-vs-inference-effect decomposition in `#der-class-coercion-via-wrapping` §Discussion “Two senses of component competence.”
 - **Open: behavioral-monitor construction.** Past the boundary only the behavioral bound $I(A(q_M); G^{\text{op}} \mid q_M)$ remains; constructing a consistent estimator for it (and for the latent G^{op} it conditions on) is the same open follow-on as the $\hat{\kappa}_{\text{processing}}$ estimator of `#der-directed-separation` and the decomposed-estimator spike candidate recorded in `#der-class-coercion-via-wrapping` Working Notes.
-

Appendix AE

Derived: Belief-Strategy Attractors Under Σ -Source Coupling

A Class 2 agent whose belief-update coupling is sourced from the strategy Σ_t — rather than from the objective O_t — closes a feedback loop $M_t \rightarrow \Sigma_t \rightarrow f_M \rightarrow M_t$ through the orient cascade. The loop admits *self-stabilizing fixed points* (M^*, Σ^*) in which the agent’s belief stays misaligned with the environment indefinitely: evidence arrives, but a strategy committed to a particular belief depresses the effective belief-update gain on evidence that would refute that strategy, locking (M, Σ) in a joint attractor. O_t -source coupling under the same structural form does *not* produce such attractors — because O_t is exogenous to M_t in steady state per the orient cascade, the feedback loop $M \rightarrow O \rightarrow f_M \rightarrow M$ does not close in nominal operation.

The result is the formal statement of the long-recognized empirical asymmetry between *sunk-cost commitment cascades* (self-reinforcing) and *identity-driven motivated reasoning* (biased but not divergently runaway): the cascade-vs-bias structural distinction is forced by the orient cascade’s exogeneity of O and not-exogeneity of Σ . The asymmetry holds *robust qualitatively* across the general nonlinear $K(\Sigma)$ form and *exact* under the linearized fixed-point analysis with a posited multiplicative gain. The result strengthens to *exact* under a derivation of $K(\Sigma)$ from utility-cost-of-belief-revision (a recommended follow-on sub-spike).

The derivation matches and extends the third row of #disc-adversarial-coupling-pressure’s mechanism table (sunk-cost engineering as $\Sigma_t \rightarrow M_t$). The mechanism table named the channel; this segment derives the dynamical consequence — that channel produces self-stabilizing fixed points, not merely bias.

Formal Expression.

Setup. A Class 2 agent has Σ_t -source coupling at stage P3 (aggregation) per #disc-partial-coupling-pathways. The coupled aggregation takes multiplicative process-form:

Formulation (strategy-source-coupled-aggregation).

$$M_{\tau+} = M_{\tau-} + K(\Sigma_{\tau-}) \cdot \ell_{\tau} \quad (\text{AE.1})$$

where ℓ_{τ} is the likelihood signal from upstream stages (goal-blind by assumption — only P3 couples; if upstream stages also couple, the cascade-propagation result of #disc-partial-coupling-pathways applies and the analysis localizes at the upstream-most coupled stage), and $K : \mathcal{S} \rightarrow \mathbb{R}^+$ is a Σ -dependent positive gain.

Sunk-cost is the canonical instantiation: $K(\Sigma)$ decreases as Σ commits more heavily to specific beliefs about reality. Formally — as the projection of Σ onto strategies presuming $M = m^*$ grows, K falls for evidence challenging m^* . The mechanism is utility-cost driven: the cost of revising M_t is bundled with the cost of abandoning Σ_t (the agent would lose the sunk strategic investment if it abandons the strategy), and the agent's belief-update gain effectively reflects this bundled cost. A derivation of $K(\Sigma)$ from a utility-cost analysis of belief-revision-under-strategic-commitment is a recommended sub-spike (see Working Notes).

The strategy is updated by the orient cascade per #der-orient-cascade:

Definition (orient-cascade-Sigma-update).

$$\Sigma_{\tau+} = f_{\Sigma}(M_{\tau+}, O_t, \Sigma_{\tau-}). \quad (\text{AE.2})$$

Closed-loop dynamics.

Substituting,

Derived (closed-loop-coupling, from #form-strategy-source-coupled-aggregation + #der-orient-cascade).

$$M_{\tau+} = M_{\tau-} + K(\Sigma_{\tau-}) \cdot \ell_{\tau}, \quad \Sigma_{\tau+} = f_{\Sigma}(M_{\tau-} + K(\Sigma_{\tau-})\ell_{\tau}, O_t, \Sigma_{\tau-}). \quad (\text{AE.3})$$

This is a closed-loop system in (M, Σ) with M_t entering $\Sigma_{\tau+}$ through the orient cascade and $\Sigma_{\tau-}$ entering $M_{\tau+}$ through the coupled aggregation.

Fixed-point structure.

A fixed point (M^*, Σ^*) satisfies

Derived (fixed-point-conditions).

$$K(\Sigma^*) \cdot \ell^* = 0, \quad \Sigma^* = f_{\Sigma}(M^*, O_t, \Sigma^*). \quad (\text{AE.4})$$

The condition $K(\Sigma^*) \cdot \ell^* = 0$ admits two qualitatively distinct solutions:

(a) Honest convergence: $\ell^* \rightarrow 0$ — no further evidence to integrate. M^* is the agent's converged belief; the asymptotic dynamics match standard Bayesian convergence under a stationary evidence stream.

(b) Pathological attractor: $K(\Sigma^*) \rightarrow 0$ — the gain has collapsed *even though* $\ell^* \neq 0$. Evidence arrives that would, under honest gain, update M — but the gain has fallen near zero on the relevant subspace, so the evidence does not propagate. The agent's M stays at m^* ; Σ stays committed to the strategy that presumes m^* ; the system rests in a (m^*, Σ^*) attractor misaligned with the environment.

Case (b) is the **belief-strategy attractor** — the formal structure of sunk-cost commitment cascades.

Linearized stability.

Linearize the closed-loop dynamics around (M^*, Σ^*) . Define perturbations $\delta M_\tau = M_\tau - M^*$, $\delta \Sigma_\tau = \Sigma_\tau - \Sigma^*$, and partial derivatives at the fixed point:

Derived (linearized-stability), exact under linearization.

$$K^* = K(\Sigma^*), \quad K'^* = \partial_\Sigma K|_{\Sigma^*}, \quad A = \partial_M f_\Sigma|_{(M^*, \Sigma^*)}, \quad B = \partial_\Sigma f_\Sigma|_{(M^*, \Sigma^*)}. \quad (\text{AE.5})$$

Treating ℓ as a function of M for the locally-linearized evidence-arrival process — $\ell \approx -L\delta M$ for some positive-semi-definite information matrix L — the linearized dynamics are:

$$\delta M_{\tau+} = (I - K^*L)\delta M_{\tau-} + (K'^*\ell^*)\delta \Sigma_{\tau-}, \quad (\text{AE.6})$$

$$\delta \Sigma_{\tau+} = A\delta M_{\tau+} + B\delta \Sigma_{\tau-}. \quad (\text{AE.7})$$

The Jacobian of the closed-loop map:

$$J = \begin{pmatrix} I - K^*L & K'^*\ell^* \\ A(I - K^*L) & AK'^*\ell^* + B \end{pmatrix}. \quad (\text{AE.8})$$

The fixed point (M^*, Σ^*) is a stable attractor iff all eigenvalues of J have modulus less than 1. In the sunk-cost-collapse regime $K^* \rightarrow 0$, the top-left block $I - K^*L \rightarrow I$ — there is *no contraction in the M direction* — and stability is determined by the Σ -direction block B . If B has spectral radius less than 1 (the orient cascade is contractive in Σ near Σ^* , which is typical for strategy-update dynamics with friction), the joint fixed point is stable *despite the absence of M -direction contraction* — the strategy stabilizes around its commitment, and the absence of belief-update propagation keeps M near m^* .

Result (pathological-attractor-condition), conditional.

This is the formal statement: **Σ -source coupling with multiplicative process-form admits self-stabilizing attractors in which the agent's belief remains misaligned with the environment indefinitely.**

Contrast — O_t -source coupling produces no such attractor.

Substitute the same coupling form with O as source:

Derived (O -source-no-attractor), conditional on orient-cascade exogeneity.

$$M_{\tau+} = M_{\tau-} + K(O_t) \cdot \ell_\tau. \quad (\text{AE.9})$$

Per #der-orient-cascade, O_t revises only when forced — under nominal operation, O_t is *exogenous* to M_t in steady state. In steady state $O_t = O^*$ is constant; $K(O_t) = K(O^*)$ is constant. The closed-loop dynamics reduce to a standard contractive Bayesian update with a goal-shaped (but fixed) gain:

$$\delta M_{\tau+} = (I - K(O^*)L)\delta M_{\tau-} + 0 \cdot \delta O_{\tau-}. \quad (\text{AE.10})$$

The O -feedback term vanishes because O does not update from M in steady state. The map is contractive whenever $K(O^*)L$ has appropriate properties; the agent's belief has a *bias* (the gain $K(O^*)$ may be smaller than the honest gain,

producing slower convergence or a non-zero asymptotic residual relative to the true environment) but it has no feedback-induced *foreclosure* of belief-revision.

Pure O_t -source coupling produces bias but not runaway commitment. The source asymmetry is structural: it follows directly from the orient cascade's exogeneity of O and not-exogeneity of Σ , not from any additional property of K .

Source-asymmetry result.

Pure Σ_t -source coupling at any stage of f_M closes a feedback loop $M \rightarrow \Sigma \rightarrow f_M \rightarrow M$ through the orient cascade and admits self-stabilizing belief-strategy attractors. Pure O_t -source coupling at the corresponding stages does not, because the orient cascade's exogeneity of O in steady state prevents the loop from closing.

Derived (source-asymmetry-attractor), robust qualitative (exact under linearization with multiplicative K).

The asymmetry is *structural* — it follows from the topology of the orient cascade, not from a parametric difference between O and Σ . Any mixed-source coupling $R = \{O, \Sigma\}$ inherits the Σ -component's attractor possibility — even moderate Σ -source coupling, paired with strong O -source coupling, can produce the attractor through the Σ channel alone.

Epistemic Status.

Conditional. Max attainable under current derivation: *exact* for the linearized fixed-point analysis under the posited multiplicative gain $K(\Sigma)$ and the orient-cascade exogeneity of O ; *robust qualitative* for the general nonlinear K and the broader claim that any closed-loop topology $M \rightarrow \Sigma \rightarrow f_M \rightarrow M$ admits the attractor structure.

Stated conditions:

(R1) **Multiplicative gain form.** $K(\Sigma)$ is posited as a positive scalar (or pointwise-positive matrix) gain modulating the aggregation step's update magnitude. The form is structurally motivated by the empirical mechanism (sunk-cost bundles the cost of revising M with the cost of abandoning Σ , which in linearization presents as gain suppression) but not derived from a more primitive AAT mechanism. A derivation of $K(\Sigma)$ from utility-cost analysis of belief-revision-under-strategic-commitment is the gating sub-spike for tier-upgrade to *exact*.

(R2) **Orient-cascade exogeneity of O in steady state.** O revises only when forced per #der-orient-cascade. Agents that violate this — continuously updating O as a function of M — collapse the source asymmetry: O -coupling then closes a feedback loop too, but the orient cascade explicitly forbids this in canonical AAT. Identity-binding under sustained adversarial pressure may be the operational regime where O effectively updates from M (per #disc-adversarial-coupling-pressure); this is where the asymmetry attenuates.

(R3) **Linearization validity.** The fixed-point stability analysis is local. Generalization to nonlinear K , f_Σ preserves the *qualitative* asymmetry (the feedback-loop topology is independent of linearization) but the explicit fixed-point structure may admit multiple attractors, limit cycles, or chaos.

(R4) **Stage of coupling.** Stated for P3 (aggregation) for definiteness; the asymmetry is *generic across stages* — the closed-loop topology depends on which source is coupled, not on which pipeline stage hosts the coupling. P3 is the canonical sunk-cost site; identity-protective consolidation (P4 / O) and frame-coupling (P1 / Σ) instantiate the asymmetry at other stages.

Discussion.

The structural reason sunk-cost has a different empirical signature than motivated reasoning. Sunk-cost cascades are documented as *self-reinforcing* in psychology (Staw 1976; Brockner 1992; Arkes & Blumer 1985; Sleesman et al. 2012 meta-analysis) and in organizational behavior. Identity-driven motivated reasoning is documented as producing *bias* but not divergent attractors except in identity-bound communities reinforcing priors. The asymmetry across these literatures has not, to our knowledge, been derived as a structural theorem; the empirical clustering is well-known but the structural reason — strategy is endogenous to belief via the cascade, while objective is exogenous in steady state — is the AAT-internal contribution. The derivation above formalizes the empirical-clustering observation as a consequence of the orient cascade’s topology.

Implication for adversarial pressure. #disc-adversarial-coupling-pressure’s sunk-cost-engineering mechanism is the adversary engineering precisely this structure: by inducing the target to commit publicly or expensively to a strategy that presumes a particular belief, the adversary locks the target’s (M, Σ) into the attractor regime. Once the attractor is established, the adversary can stop driving and the target’s own internal dynamics maintain the misaligned belief. This is the structural reason sunk-cost engineering is so effective in propaganda, influence operations, and social engineering: the target’s own machinery does the maintenance work after the initial commitment.

Implication for class-coercion-via-wrapping. A Class 3 (Coupled) component wrapped to Class 1 status via the construction of #der-class-coercion-via-wrapping requires extra care when the un-wrapped component has Σ -source coupling. Even W_1 (strict wrapping — no G in the belief-update query) may be insufficient if the component’s *internal* Σ is influenced by historical query content that carries strategic context, and that internal Σ then suppresses the gain on subsequent goal-blind queries. The corrective is a Σ -channel-suppressed W_1 : the wrapper must additionally hold the component’s strategic context fixed across calls or strip strategic-context-bearing content from the queries. This is a finer-grained design constraint than the wrapping-construction segment currently surfaces; it follows directly from the source-asymmetry result.

Implication for compositional dynamics. The composite-level class inheritance table in #der-directed-separation tracks composite Class membership but not composite sub-type. A composite of Class 1 (Separated) sub-agents whose routing is goal-dependent (R_t depends on G_t^c) is Class 3 at the composite level per the inheritance table — but the source-asymmetry result of this segment further refines the composite’s sub-type: if the routing’s G^c -dependence is on the composite’s strategy Σ^c rather than its objective O^c , the composite admits belief-strategy attractors at the composite level. This is the operational reason organizations with strong strategic commitments and goal-dependent internal routing exhibit sunk-cost dynamics at the institutional scale.

The diagnostic question for a Class 2 agent. When inspecting a Class 2 agent, the diagnostic question that distinguishes attractor-vulnerable from bias-vulnerable sub-types is: *is the coupling sourced from O_t or from Σ_t ?* The question is sharper than asking the aggregate $\kappa_{\text{processing}}$ magnitude, because the topological consequence (attractor possibility) depends on the source, not on the magnitude. Empirically, the source can be probed by varying O_t at fixed Σ_t vs Σ_t at fixed O_t and observing which probe condition reveals the gain modulation. In setups where holding Σ fixed is operationally difficult, the diagnostic falls back to the dynamical-signature analysis: does the agent’s belief, under sustained counter-evidence, *converge slowly with bias* (O -source signature) or *fail to converge at all* in a self-stabilizing pattern (Σ -source signature)?

Findings.

Belief-Strategy Attractors From Σ -Source Coupling (Source Asymmetry). **Brief:** An agent whose belief-update is shaped by *what it currently wants to be true* will be biased — its beliefs will lean toward the goal — but a thoughtful third party can in principle measure the bias and correct it, and the agent’s beliefs will eventually catch up to reality as evidence accumulates. An agent whose belief-update is shaped by *what plan it is currently committed to* is structurally different: the strategy commitment depresses the agent’s belief-update gain on evidence that would invalidate the strategy, the depressed gain keeps belief at the strategy-consistent point, and the unchanging belief sustains the strategy

commitment. The agent locks into a joint *belief-and-strategy attractor* — a fixed point where the agent stays misaligned with reality indefinitely, with no internal mechanism to escape. The structural reason: the agent's strategy is updated by the agent's beliefs, so coupling beliefs to strategy closes a feedback loop. Beliefs and objectives are not coupled the same way — objectives are revised only when forced, so coupling beliefs to objectives doesn't close the loop. This is the structural reason sunk-cost cascades self-reinforce in a way that identity-driven motivated reasoning does not.

Impact: Names the structural origin of the sunk-cost / motivated-reasoning empirical asymmetry — both have been documented across psychology and organizational behavior, but the structural reason (orient cascade endogeneity of Σ vs exogeneity of O) has not, to our knowledge, been stated as a derived theorem. Sharpens the wrapping-regime selection of $\#der$ -class-coercion-via-wrapping for Class 2 components with Σ -source coupling: Σ -channel-suppressed W_1 may be required even when ordinary W_1 would suffice for O -source coupling. Connects $\#disc$ -adversarial-coupling-pressure's sunk-cost-engineering mechanism (third row of the mechanism table) to a dynamical-attractor structure rather than just a channel; explains why the mechanism is so effective in propaganda and influence operations (the target's own machinery does the maintenance work). Provides a diagnostic question for inspecting Class 2 agents — source of coupling — that is sharper than the aggregate $\kappa_{processing}$ magnitude.

Novelty Claim: *Claim recognition* of the structural source asymmetry as a direct consequence of the orient cascade's topology — strategy is endogenous to belief, objective is exogenous in steady state. *Claim differentiation* on the resulting fixed-point analysis: linearized stability with $K^* \rightarrow 0$ produces an attractor in the closed-loop dynamics under the Σ -source case but not under the O -source case. The empirical clustering across the sunk-cost and motivated-reasoning literatures is well-established; the AAT-internal contribution is naming the structural reason and stating it as a derived theorem.

Related Work:

- Staw 1976 “Knee-Deep in the Big Muddy” *Organizational Behavior and Human Performance* (published 1976, intuition-only date) — *empirical instantiation supporting* — the canonical psychological study of self-reinforcing commitment to a chosen course of action; the empirical phenomenon the source-asymmetry result structurally explains.
- Brockner 1992 *Academy of Management Review* “The Escalation of Commitment to a Failing Course of Action” — *empirical instantiation supporting* — synthesizes the sunk-cost / commitment-cascade literature; same empirical anchor.
- Arkes & Blumer 1985 *Organizational Behavior and Human Decision Processes* “The Psychology of Sunk Cost” — *empirical instantiation supporting* — foundational behavioral-economics treatment of sunk-cost effects.
- Sleesman et al. 2012 *Academy of Management Journal* meta-analysis on escalation of commitment — *empirical instantiation supporting* — quantitative synthesis of the sunk-cost-cascade literature; empirical confirmation of the self-reinforcement signature.
- Kunda 1990 *Psychological Bulletin* “The Case for Motivated Reasoning” — *empirical instantiation supporting (contrast case)* — the canonical motivated-reasoning treatment; produces bias but not the divergent-attractor signature of sunk-cost — the empirical contrast the source-asymmetry result structurally explains.
- Kahan cultural-cognition project (various years) — *empirical instantiation supporting (contrast case)* — identity-driven motivated reasoning produces persistent bias but typically not runaway divergence except when paired with composite-level identity-binding community dynamics; the contrast with sunk-cost is exactly the source-asymmetry distinction.
- Bénabou & Tirole 2002, 2011 *Quarterly Journal of Economics* on identity economics and motivated beliefs — *adjacent literature, not yet searched* — formal models trading off truth vs identity; possibly anticipates aspects of the source distinction.

Search Log:

- 2026-05-22 (*intuition-only*): The empirical sunk-cost / commitment-cascade literature is well-established and well-known to converge on the self-reinforcement signature; the empirical motivated-reasoning literature is similarly well-established and converges on the bias-without-runaway signature. The *structural derivation* of the asymmetry from the orient cascade's topology has not been searched. Recommended Pillar-style targeted search before tier upgrade — primary targets are the formal-models side of behavioral economics (Bénabou-Tirole identity economics; Akerlof-Kranton identity economics) and the formal-models side of motivated-reasoning psychology (Kruglanski lay epistemics; possibly Shah-Friedman-Kruglanski goal-shielding 2002 *JPSP*). The active-inference / control-as-inference literatures are unlikely to have this specific result (they don't typically separate O and Σ as distinct goal-state objects), but should be confirmed.

Working Notes.

- **Tier-upgrade gating sub-spike — derivation of $K(\Sigma)$ from utility-cost analysis.** The multiplicative-gain form $K(\Sigma)$ is *posited* per condition (R1). The empirical mechanism (sunk-cost bundles the cost of revising M with the cost of abandoning Σ) suggests the effective dynamics are gain-suppressing in the linearized regime, but the explicit derivation has not been produced. Sub-spike scope: model the agent's belief-update as an optimization minimizing $\text{prediction-error}(M, e) + \lambda \cdot \text{strategy-revision-cost}(\Sigma, M)$, derive the effective belief-update first-order condition, and verify that the linearized form is gain-modulation with $K \propto 1/(1 + \lambda \cdot \partial_M^2 \text{strategy-revision-cost})$. If this matches the multiplicative form, the result strengthens to *exact* under the linearization. If the linearization yields a different form (e.g., additive regularization rather than multiplicative gain), the result is refined accordingly. Recommended as the higher-priority of the two disc-partial-coupling-pathways-gated sub-spikes.
- **Tier-upgrade gating sub-spike — Pillar prior-art search.** The empirical clustering across sunk-cost and motivated-reasoning literatures is well-established; the structural derivation as a theorem is plausibly novel under search-depth conducted. Recommended Pillar-style targeted search across behavioral economics (Bénabou-Tirole, Akerlof-Kranton), motivated-cognition psychology (Kruglanski, Shah-Friedman-Kruglanski), and self-deception philosophy (Mele, von Hippel-Trivers). See #disc-partial-coupling-pathways Working Notes for the broader scope.
- **Composite-level extension.** A composite of Class 1 sub-agents with goal-dependent routing where the goal-dependence is on the composite's *strategy* Σ^c (rather than its objective O^c) inherits the attractor possibility at the composite level. This refines the composite-level class inheritance table in #der-directed-separation. Extending the table to track composite sub-type (and hence composite attractor-vulnerability) is the structural-composition follow-on noted in #disc-partial-coupling-pathways Working Notes.
- **Empirical estimator — distinguishing O -source bias from Σ -source attractor under behavioral observation.** Conceptually distinct (slow biased convergence vs convergence-failure); in practice may require careful experimental design. The dynamical signature is the cleanest discriminator: under sustained counter-evidence, O -source coupling produces convergence asymptotically (with bias-magnitude proportional to the gain reduction); Σ -source coupling produces non-convergence (the gain falls toward zero faster than the evidence accumulates). The empirical literatures on sunk-cost vs motivated-reasoning have already documented this clustering; a behavioral-estimator construction that decisively distinguishes the two on a single agent under controlled probing is a separate work item.
- **Connection to #deriv-strategic-composition.** The strategic-composition machinery handles composite-level dynamics under partially-opposing objectives, with the dynamic-regime axis (#disc-dynamic-regime-axis) carrying the goal-alignment effects independently of architectural class. The source-asymmetry result of this segment is at the *single-agent* level — it constrains the structural properties of an individual Class 2 agent. The composite-level extension (above) would bridge the two; not pursued in this segment.
- **Brief is below Feynman criterion.** The current Brief states the result by contrasting two empirically-familiar phenomena (sunk-cost cascade vs identity-driven motivated reasoning), which is closer to recognition than to an isomorphic physical analog. A candidate analog reaching for the Feynman criterion: imagine a thermostat whose *current setpoint* (objective) is fixed exogenously by a building operator, vs one whose setpoint *itself drifts* based on the room temperature it has been recently

reading — the second creates a feedback loop where the thermostat slowly forgets what temperature the room is *supposed* to be at, because the goalpost has moved with the observation. The first thermostat reaches the setpoint with some bias if the temperature sensor is off; the second can drift indefinitely. The thermostat analog reaches for the closed-loop topology but doesn't capture the multiplicative gain mechanism — the load-bearing structure is in *which way* the loop closes, not just that one closes. The analog needs sharpening.

Appendix AF

Discussion: The Dynamic-Regime Axis — Contraction / Equilibrium / Cyclic-Distributional / Mean-Field

AAT composites differ on at least three independent structural axes: their *architectural class* (Class 1 Separated / Class 2 Partial / Class 3 Coupled, per #der-directed-separation) — about processing topology; their *dynamic regime* (this segment) — about the structural type of the joint-dynamics fixed point and the Lyapunov certificate available; and their *macro-state type* (state-variable / fixed-point object / distributional) — about what kind of object the composite's G_t^c is. This segment names the dynamic-regime axis as a first-class structural classifier, in vocabulary the framework already runs implicitly across #deriv-strategic-composition's α'/β' partition, #form-composition-closure's contraction presupposition, #disc-separability-pattern's Contraction ladder, #scope-composite-agent's scope-route disjunction, and #impl-strategic-composition's “contraction-to-equilibrium hand-off” framing.

The four-tier ladder is recognized prior art. This is integration, not invention. The partition AAT names — contraction-regime / equilibrium-regime / cyclic-distributional-regime / mean-field-equilibrium-regime — appears (under different vocabulary, with isomorphic structural content) across at least five neighboring literatures. Candogan, Menache, Ozdaglar & Parrilo's 2011 Hodge-style decomposition of finite games into potential + harmonic + nonstrategic components supplies a coordinate-free organizing structure for the equilibrium-regime / cyclic-regime boundary cleaner than any tier-wise enumeration; Letcher, Balduzzi et al.'s 2019 *Differentiable Game Mechanics* lifts the same potential/Hamiltonian decomposition to the dynamics side; Sandholm's 2010 *Population Games and Evolutionary Dynamics* organizes the potential/stable/cyclic taxonomy at the population level; Hofbauer & Sigmund's 1998 *Evolutionary Games and Population Dynamics* names the phase-portrait classification (equilibrium vs cycle vs heteroclinic) directly; Omidshafiei et al.'s 2019 α -Rank construction uses cyclic-dynamics-as-meaning as a solution concept; and Papadimitriou-Piliouras's 2018 *From Nash Equilibria to Chain Recurrent Sets* establishes the chain-recurrent set as the natural object when dynamics fail to converge. AAT's contribution is *not* the classification itself; it is (a) the consolidation of the partition into AAT-internal vocabulary that travels with #scope-composite-agent's scope-route disjunction and #disc-separability-pattern's separable-core / structured-repair / general-open shape, (b) three derivable structural identities that name what the axis carries within an integrated agent theory, and (c) cross-axis observations about how the axis composes with architectural class and macro-state type.

This segment is the *positive-half companion* to the broadened #disc-identifiability-floor Instance 3, which names the negative half — regime identification from component marginals is structurally impossible without composite-level coupling-topology observation. The two together mark AAT's scope on the axis at both extremes, parallel to how #disc-separability-pattern and #disc-identifiability-floor jointly mark scope on the identification axis.

The pattern.

Each composite agent satisfying #scope-composite-agent admits a dynamic regime — a structural property of its joint best-response dynamics that is *determined by the game structure* (the sub-agent objectives, the strategy spaces, the payoff curvatures) rather than by any particular trajectory the dynamics realize. The regime determines the strongest Lyapunov-class certificate the dynamics admit relative to their fixed-point set, and that certificate determines the operational machinery available — exponential convergence near a unique attractor, exponential convergence near an equilibrium set, polynomial convergence in distribution, or measure-space dynamics at the population-scale limit.

The classification is *not behavioral* (it does not depend on observing trajectories) and it is *not parametric* (it does not depend on a continuous tuning parameter). It is structural and discrete, the same shape #der-directed-separation's architectural classification has — a partition determined by the game structure, with sub-tier classifications within each regime.

The four tiers.

Following the order of progressively weakening Lyapunov certificate:

Table AF.1

<i>Tier</i>	<i>Fixed-point structure</i>	<i>Lyapunov certificate</i>	<i>Macro-state type</i>	<i>Convergence rate</i>	<i>Scope route entry</i>
R0 Contraction	Unique attracting $X^* \in \mathcal{X}^c$	Joint Lyapunov V , global or basin-restricted	State-variable X^*	Exponential at α_{joint}	C-i / C-ii / C-iii
R1 Equilibrium (α')	Fixed-point set $\mathcal{E} \subset \mathcal{X}^c$, locally attracting	Potential Φ (Monderer-Shapley) or weighted-norm (Rosen)	Fixed-point object $\mathcal{E}$	Exponential near $\mathcal{E}$ at $\alpha_{\text{joint}}^{(j)}$	C-iv + potential/monotone
R2 Cyclic-distributional (β')	No pure-strategy fixed point, or saddle-only	None on $\mathcal{X}^c$; doubling-trick CCE on $\Delta(\mathcal{X}^c)$	Distributional object $\mu \in \Delta(\mathcal{X}^c)$	Polynomial $O(1/\sqrt{T})$	C-iv without potential/monotone
R3 Mean-field equilibrium	Distributional fixed point in measure space	MFG-specific (Lasry-Lions monotonicity)	Distribution-on-population $\rho \in \mathcal{P}(\mathcal{A})$	MFG-specific, often exponential	Population-scope gap

Each tier inherits the tier-internal classification of its closest existing AAT machinery — R0's Tier 1 / Tier 2 / Tier 3 sub-classification from #disc-separability-pattern's Contraction ladder; R1's potential / monotone / general sub-tier per #deriv-strategic-composition's sub-scope α' ; R2's regret-bound sub-classification per Hart-Mas-Colell 2000; R3's MFG-uniqueness sub-classification per Lasry-Lions 2007.

R0 — Contraction-regime. The joint dynamics admit a unique attracting fixed point $X^* \in \mathcal{X}^c$ and a Lyapunov function $V : \mathcal{X}^c \rightarrow \mathbb{R}_{\geq 0}$ with $V(X^*) = 0$, $V(X) > 0$ for $X \neq X^*$, such that

$$\frac{d}{dt} V(X_t^c - X^*) \leq -\alpha_{\text{joint}} \cdot d(X_t^c, X^*)^2 \quad \forall X_t^c \in \mathcal{B}(X^*) \quad (\text{AF.1})$$

for some basin $\mathcal{B}(X^*)$ and contraction rate $\alpha_{\text{joint}} > 0$. Macro-state $X^* \in \mathcal{X}^c$ is a state-variable. Composites satisfying scope routes C-i (shared composite objective), C-ii (hierarchical derivation), or C-iii (mutual-benefit alignment) under matched-Tier sub-agents per `#deriv-critical-mass-composition` land here. The classical AAT machinery applies: `#result-sector-persistence-template`, `#form-composition-closure` bridge lemma, `#der-team-persistence`, `#der-tempo-composition`.

R1 — Equilibrium-regime (α'). The joint dynamics admit a fixed-point set $\mathcal{E} = \{X_j^*\}_{j \in J}$ with $|J| \geq 1$ where each X_j^* is a Nash equilibrium. Lyapunov-on-deviation holds locally near each equilibrium under either potential-game structure (Monderer-Shapley 1996) or monotone-game structure (Rosen 1965 diagonal strict concavity). The macro-state is a fixed-point object — either a single Nash X^* if unique, or the equilibrium set if multiple. Composites satisfying scope route C-iv with potential/monotone structure land here; `#deriv-strategic-composition`’s Cournot worked example is the canonical instantiation. The sector-persistence template transfers with the joint potential’s curvature (potential case) or the joint Jacobian’s symmetric part (monotone case) supplying $\alpha_{\text{joint}}^{(j)}$ at each equilibrium.

R2 — Cyclic-distributional-regime (β'). The joint dynamics admit no fixed point in pure-strategy space, or fixed points that are not local attractors (saddle-Nash in zero-sum games where best-response orbits without approaching). No-regret dynamics drive the empirical joint distribution to the coarse correlated equilibrium set (Hart-Mas-Colell 2000):

$$\frac{1}{T} \sum_{t=1}^T \mathbb{1}\{\pi_t^c \in \text{CCE}\} \rightarrow 1 \quad \text{at rate } O(1/\sqrt{T}). \quad (\text{AF.2})$$

The macro-state is a distributional object — a measure $\mu \in \Delta(\mathcal{X}^c)$ supported on the CCE set. Composites satisfying scope route C-iv without potential or monotone structure land here. AAT in R2 predicts only set-convergence of empirical play to CCE; it cannot predict short-run trajectory, per-sub-agent mismatch convergence, or equilibrium selection. This is the framework’s honest β' scope limit — a regime limit, not an architectural limit, distinct from any architectural-class restriction.

Cross-layer linearization fingerprint at saddle-only Nash. R2’s saddle-only Nash equilibria have a single-agent linearization-spectrum fingerprint of imaginary-axis-semisimple — the *R0-loss* rung of `#result-certificate-existence`’s single-agent certificate-strength ladder. The two R-letter ladders are layered, not collided: this segment’s R0/R1/R2/R3 classifies composite joint-dynamics on $\mathcal{X}^c$, while `#result-certificate-existence`’s R0-loss/R0-strict/R1/R2 classifies single-agent certificate-strength on the linearized error space $\mathbb{R}^n$. The bridge is via composition (the joint Jacobian at R2’s saddle-Nash linearizes to the antisymmetric R0-loss-at- $\mathcal{M} = I$ sub-case); the letters are deliberately reused to mark the bridge, not by collision.

R3 — Mean-field-equilibrium-regime. For $N \rightarrow \infty$, each agent interacts with the *distribution* of other agents’ strategies rather than with named others. The macro-state becomes a distributional fixed point in measure space; the dynamics may contract to this fixed point under explicit conditions (Lasry-Lions 2007; Huang-Malhamé-Caines 2006). R3 is distinct from R2: R2’s macro-state is a distribution-on-strategies arising as the empirical play of finite- N dynamics with saddle-or-cyclic fixed-point structure on the pure-strategy space; R3’s macro-state is a distribution-on-population-strategies arising as the thermodynamic limit with genuine attractor on the measure space. R3 is the only tier where macro-state type genuinely diverges from dynamic regime in a way the framework’s existing α'/β' does not capture. R3 is the population-scope gap that Part III currently flags.

AAT-specific derivable content.

The axis is not only consolidation. Three derivable contents fall out once it is stated formally; one is the candidate floor instance landing concurrently in `#disc-identifiability-floor` Instance 3 (broadened).

Identity: macro-state-type and dynamic regime coincide for finite N . For finite- N composites (tiers R0/R1/R2), the macro-state type is structurally identified by the dynamic regime:

- $\mathcal{R}(X^c) = R0 \Leftrightarrow G_*^c \in \mathcal{X}^c$ (state-variable).
- $\mathcal{R}(X^c) = R1 \Leftrightarrow G_*^c = \mathcal{E} \subset \mathcal{X}^c$ (fixed-point object).
- $\mathcal{R}(X^c) = R2 \Leftrightarrow G_*^c = \mu \in \Delta(\mathcal{X}^c)$ (distributional object).

Each direction follows from the entry conditions in §”The four tiers.” R0 demands unique attracting fixed point in $\mathcal{X}^c$, which gives a state-variable. R1 demands a fixed-point set under Lyapunov-on-deviation, which is a fixed-point object. R2 demands no pure-strategy fixed point plus no-regret CCE convergence, which gives a distributional object. The converse direction follows from the macro-state type forcing the available Lyapunov machinery: state-variable admits scalar Lyapunov (R0); fixed-point set admits potential or weighted-norm Lyapunov (R1); distributional macro-state admits only doubling-trick CCE convergence (R2).

R3 breaks the identity. Mean-field-equilibrium has distributional macro-state with contracting dynamics, distinct from R2’s polynomial CCE convergence. The two axes (dynamic regime, macro-state type) are genuinely distinct and the breaking case is exactly the population-scope frontier the framework currently flags as gap. The identity is the cleanest formal content the axis carries: it is the axis that unifies fixed-point structure, Lyapunov certificate, and macro-state type into a single classification for finite- N composites, with the mean-field limit as the explicit boundary case.

Transition asymmetry: descent is the default, ascent costs more. Five non-trivial transitions exist between adjacent tiers, with a structural asymmetry between descent ($R0 \rightarrow R1$; $R1 \rightarrow R2$) and ascent ($R1 \rightarrow R0$; $R2 \rightarrow R1$). Descent under generic perturbation of the objective structure is the default — alignment is fragile, potential-structure is fragile, both can be destroyed by local perturbations. Ascent requires explicit structural commitment that costs more than the descent did to acquire. Symbiogenic absorption (#hyp-symbiogenic-composition) is the structural mechanism for $R1 \rightarrow R0$ ascent; mechanism design under explicit potential-game constraints is the mechanism for $R2 \rightarrow R1$ ascent — both impose structural commitments that out-of-band negotiation, designer intervention, or institutional shaping must supply.

This is an energy-asymmetry on the dynamic-regime axis — descent is downhill, ascent is uphill, and the asymmetry is structural rather than parametric. Adversaries exploit it as the modularity-state version exploits the modularity descent (see #disc-adversarial-coupling-pressure for the modularity analog).

The transition asymmetry is currently argumentation-grade; a fully-derived statement tied to a structural-distance metric on game-rewiring space would upgrade it to derived, and this segment names the shape pending that derivation. The mechanism-design impossibility theorems (Gibbard-Satterthwaite 1973–75, Myerson-Satterthwaite 1983, Arrow 1951) are *floors on regime-ascent* under stated designer constraints; the queued #disc-implementation-impossibility meta-segment carries the formal treatment of mechanism-design impossibility as designer-side constructive-impossibility.

Lyapunov-machinery transfer $R0 \rightarrow R1$. The persistence template #result-sector-persistence-template lifts cleanly across $R0 \rightarrow R1$, with the Lyapunov function changing form at each tier:

$$R0: V_{R0}(X^c) = \|X^c - X^*\|^2 \quad \dot{V}_{R0} \leq -\alpha_{\text{joint}} V_{R0} \quad (\text{AF.3})$$

$$R1 \text{ potential: } V_{R1,\Phi}(\pi^c) = \Phi(\pi^c) - \Phi(\pi_*^c) \quad \dot{V}_{R1,\Phi} \geq \alpha_{\text{joint}} \|\nabla \Phi\|^2 \quad (\text{AF.4})$$

$$R1 \text{ monotone: } V_{R1,F}(\pi^c) = \frac{1}{2}(\pi^c - \pi_*^c)^\top P(\pi^c - \pi_*^c) \quad \dot{V}_{R1,F} \leq -\alpha_{\text{joint}} V_{R1,F} \quad (\text{AF.5})$$

with P the weighted-norm matrix solving the Lyapunov equation for the linearized ∇F at π_*^c . At R2 the Lyapunov form fails — no state-space Lyapunov exists; the dynamics orbit equilibria rather than converge to them. The doubling-trick (Hart-Mas-Colell 2000) gives a regret bound, which translates to empirical-distribution convergence to CCE at $O(1/\sqrt{T})$ rate. The transfer makes the framework's existing α'/β' partition explicitly a regime-axis property rather than scattered across #deriv-strategic-composition.

What the axis lets the framework consolidate.

The axis already lives across ten-plus existing segments under different names. Surfacing it as a formal classifier consolidates these implicit usages into one explicit definition:

Table AF.2

<i>Existing segment / construct</i>	<i>Axis-vocabulary content</i>
#scope-composite-agent routes C-i/ii/iii	Entry conditions for R0
#scope-composite-agent route C-iv	Entry condition for R1 + R2 (potential/monotone $\rightarrow$ R1; without $\rightarrow$ R2)
#form-composition-closure contraction presupposition	R0 machinery
#result-sector-persistence-template	R0 sector constant + persistence
#deriv-critical-mass-composition matched-Tier closed form	R0 sub-tier specialization
#der-team-persistence cooperative limit	R0 limiting case
#deriv-strategic-composition sub-scope α'	R1 with potential/monotone Lyapunov transfer
#deriv-strategic-composition sub-scope β'	R2 with CCE convergence only
#impl-strategic-composition “contraction-to-equilibrium hand-off”	R0 $\rightarrow$ R1 transition
#disc-separability-pattern Contraction ladder (row 4)	R0 sub-tier classification
#hyp-symbiogenic-composition	R1 $\rightarrow$ R0 transition mechanism (structural absorption)
#der-adversarial-destabilization	R0 or R1 target under exogenous adversarial parameter; per-target regime distinction

The consolidation has two operational consequences. First, the cross-segment consistency check becomes mechanical — any segment using regime-axis vocabulary can be cross-referenced against this segment's four-tier specification. Second, the meta-segment cross-references to #disc-separability-pattern (Contraction ladder as R0 sub-tier classifier) and to broadened #disc-identifiability-floor Instance 3 (regime-identification floor) form the integration backbone of how the axis sits inside the wider meta-pattern landscape.

Cross-axis interactions.

Three structural axes coexist in the framework after the dynamic-regime axis surfaces:

- **Axis A — Architectural class.** Processing topology — Class 1 / 2 / 3 per #der-directed-separation. Determined by sub-agent class plus routing structure plus shared-substrate-with- G^c -allocation.
- **Axis B — Dynamic regime.** Joint-dynamics fixed-point structure plus Lyapunov certificate — R0 / R1 / R2 / R3. Determined by objective alignment plus game structure plus population scale.
- **Axis C — Macro-state type.** Type-theoretic shape of G_*^c — state-variable / fixed-point object / distributional. For finite N , identified by Axis B per §”Identity”; for $N \rightarrow \infty$ (R3) the identity can break.

Axes A and B are independent. The full cross-product is meaningful:

Table AF.3

	<i>R0 contraction</i>	<i>R1 equilibrium</i>	<i>R2 cyclic-distributional</i>
Class 1 Separated	Kalman + LQR team, shared target	Cournot duopoly (distinct hardware)	Rock-paper-scissors players on distinct hardware
Class 2 Partial	Biological cortex in coordinated task	Hybrid AI–biological with partial coupling, strategic interaction	Hybrid in cyclic game
Class 3 Coupled	Multi-LLM shared substrate, aligned objective (wrapping construction available)	Multi-LLM shared substrate, strategic interaction (wrapping available)	Multi-LLM shared substrate, cyclic interaction

The cells are operationally distinct. The wrapping construction (`#der-class-coercion-via-wrapping`) operates *only* on Axis A — it converts Class 3 $\rightarrow$ Class 1 architecturally at the wrapper level, at structural-leakage cost. It does *not* change Axis B regime. A wrapped strategic composite remains in R1; the wrapping recovers Class 1 architecturally but the dynamics still admit equilibrium-shaped failure modes (multi-equilibria, saddle-Nash, equilibrium selection ambiguity). Symmetrically, alignment work operates only on Axis B — re-aligning objectives moves the composite from R1 to R0, without changing architectural class. Symbiogenic absorption is the structural mechanism for the R1 $\rightarrow$ R0 transition (`#hyp-symbiogenic-composition`).

Two distinct adversarial pressures act on the two axes. `#disc-adversarial-coupling-pressure` names the adversarial pressure on Axis A — adversaries strategically drive a target’s modularity downward because coupling expands attack surface. A separate adversarial pressure — call it *strategic divergence pressure* — drives a target’s Axis B regime downward (introduces objective divergence into a previously-aligned team, moving R0 $\rightarrow$ R1; introduces cyclic structure by mechanism manipulation, moving R1 $\rightarrow$ R2). The two are structurally distinct adversarial moves with distinct defenses (architectural scaffolding for Axis A; mechanism-design commitment for Axis B). The M4 meta-segment `#disc-modularity-state-dynamics` covers the three-operation modularity-state picture on Axis A; the dynamic-regime axis here covers the parallel three-operation picture on Axis B (alignment-strengthening self-driven-ascending; strategic-divergence-acquisition self-driven-descending; adversarial-divergence-pressure externally-driven-descending).

Epistemic Status.

Discussion-grade at the meta-pattern level. The segment articulates an axis-vocabulary consolidation rather than a derivation of new machinery. The four-tier partition itself is integration of recognized prior art — Candogan-Menache-Ozdoglar-Parrilo 2011 (Hodge decomposition), Letcher-Balduzzi et al. 2019 (Hamiltonian dynamics decomposition), Sandholm 2010 (population-games taxonomy), Hofbauer-Sigmund 1998 (phase-portrait classification), Omidshafiei et al. 2019 (α -Rank), Papadimitriou-Piliouras 2018 (chain-recurrent sets as natural object).

The AAT-specific derivable contents ($\S$ “AAT-specific derivable content”) are at distinct tiers: the macro-state-type / regime identity is *exact* for finite N given the entry conditions; the transition asymmetry is *argumentation-grade* pending the J9 derivation spike with a structural-distance metric; the Lyapunov-machinery transfer R0 $\rightarrow$ R1 is *exact* by direct construction (the algebra is inherited from `#deriv-strategic-composition` sub-scope α').

Max attainable: *discussion-grade* for the meta-pattern (it is an organizing principle, not a derivation); *exact* for the macro-state-type / regime identity; *exact* for the Lyapunov transfer; *argumentation-grade* for the transition asymmetry pending the J9 derivation.

Discussion.

The cross-segment integration consequence. Once the axis is named, #impl-strategic-composition's §"Modular safety architectures fail under goal divergence" implication has a cleaner derivation than the original architectural-class-change framing. The failure mechanism is *regime change* ($R0 \rightarrow R1$ under goal divergence) rather than *architectural-class change* (which does not happen — strategic composition with goal-blind routing leaves the composite architecturally Class 1). Modular safety guarantees designed for contraction-regime composites depend on the unique-attracting fixed point and the global Lyapunov contraction; under $R0 \rightarrow R1$ transition the unique attractor becomes a set of equilibria, saddle-point Nash breaks Lyapunov contraction, and multi-equilibria break uniqueness. Constitutional AI re-teaming and mesa-optimizer formation are the empirical instantiations. The dynamic-regime axis re-derivation gives the same conclusion on stronger ground.

Mechanism design as regime-shaping. An external designer who can shape the sub-agent objectives can target a regime — bring R2 down to R1 by adding potential structure, or R1 down to R0 by adding alignment. The mechanism-design impossibility theorems (Gibbard-Satterthwaite, Myerson-Satterthwaite, Arrow) are floors on what *external designers* can achieve in the regime-ascent direction under stated incentive-compatibility and welfare constraints. The queued #disc-implementation-impossibility meta-segment carries the formal treatment of designer-side constructive-impossibility; this segment's transition-asymmetry argument cross-references it as the structural complement.

The dynamic-regime axis and #disc-modularity-state-dynamics. Both axes admit a parallel three-operation modularity-like dynamics: alignment-strengthening / strategic-divergence-acquisition / adversarial-divergence-pressure on Axis B parallel to truthification / strategic-self-coupling / adversarial-coupling-pressure on Axis A. The parallel-pattern observation is itself meta-architectural content — the framework has *two* independently-shaped axes with parallel three-operation dynamics on each. Whether this parallel structure warrants a *higher-order* meta-segment (covering both Axes A and B at once) is open; the current arrangement keeps the two axis-segments separate (this one for Axis B; #disc-modularity-state-dynamics for Axis A landed 2026-05-24) and surfaces the parallel observation in each. Now that both meta-segments are stable, evaluating whether the cross-pattern is load-bearing enough for a higher-order meta-segment is a candidate follow-on.

Sub-tier structure within each regime. Each regime has its own internal classification — R0 inherits the Contraction ladder Tier 1 / Tier 2 / Tier 3 from #disc-separability-pattern; R1 has potential / monotone / general sub-tier per #deriv-strategic-composition; R2 has CCE-convergence-rate sub-classification; R3 has MFG-uniqueness sub-classification. Whether these sub-classifications unify across regimes under a single meta-classification (e.g., "Lyapunov-certificate-strength" as one axis across R0/R1/R2/R3) is a candidate for follow-on work — see Working Notes.

Findings.

The Dynamic-Regime Axis as Cross-Cutting Classifier. **Brief:** Composites differ in how their joint dynamics behave even when their components are individually clean. Think of a team rowing a boat — everyone wants the same destination, the dynamics contract to a shared target, the team is in the *contraction-regime* (R0). Now think of firms in a market — each wants to maximize their own profit, the prices settle into a stable joint equilibrium where no firm wants to deviate, the system is in the *equilibrium-regime* (R1). Now think of rock-paper-scissors players — no stable joint move exists in pure strategies, but the long-run frequencies of moves settle into a stable distribution, the system is in the *cyclic-distributional-regime* (R2). At population scale, where the natural object is the distribution of strategies across the whole population rather than named individuals, you get the *mean-field-equilibrium-regime* (R3). The regime is structural — you can read it off the game definition without running the dynamics — and it determines what kind of guarantee the framework can give you: exponential convergence in R0/R1, polynomial-in-distribution in R2, measure-space machinery in R3. The axis is independent of the architectural class (whether the agent's processing is goal-blind or not); both axes interact, and a strategic composite of architecturally-clean agents is still R1 in dynamic regime even though it remains Class 1 in architectural class.

Impact: Names a structural classifier the framework had been running implicitly across ten-plus existing segments under different vocabulary, consolidating them into one explicit axis with a clear sub-tier structure inherited from existing tiered ladders. Provides a cleaner derivation of #impl-strategic-composition’s “modular safety architectures fail under goal divergence” implication — the mechanism is regime change ($R0 \rightarrow R1$), not architectural-class change (which does not happen under goal-blind routing). Resolves the long-standing tension in #scope-composite-agent’s Working Notes between the alignment routes (C-i/ii/iii) and the strategic-equilibrium route (C-iv): they correspond to regime-tier entry conditions on the dynamic-regime axis, not to a continuous spectrum on a single scalar. Surfaces the structural-distance metric on regime-transition cost as a derivable open question (J9 spike candidate). Provides the positive-half complementarity for the broadened #disc-identifiability-floor Instance 3 (regime identification from component marginals is structurally impossible).

Novelty Claim: *Claim integration* of multi-agent dynamics classifications from at least five neighboring literatures (game-theoretic decomposition; multi-agent learning convergence taxonomy; population games; evolutionary games; mean-field games) into an AAT-internal regime-axis vocabulary that travels with the framework’s existing scope-route disjunction (#scope-composite-agent), separable-core / structured-repair / general-open shape (#disc-separability-pattern), and persistence machinery (#result-sector-persistence-template). The four-tier partition is recognized prior art; the AAT-specific contributions are (a) the macro-state-type / regime structural identity for finite N with $R3$ as the explicit boundary case where the identity breaks (*claim recognition* + *claim integration*); (b) the regime-transition cost-asymmetry framing (descent default / ascent costly) tied to mechanism-design impossibility (*claim recognition* — full derivation deferred to J9); (c) the Lyapunov-machinery transfer $R0 \rightarrow R1$ made explicit as a cross-tier consolidation of #deriv-strategic-composition sub-scope α' (*claim integration*).

Related Work:

Table AF.4

<i>ASF concern</i>	<i>Prior-art language</i>	<i>Relationship / Positioning</i>
Game decomposition (potential / harmonic / nonstrategic)	Candogan, Menache, Ozdaglar & Parrilo 2011, “Flows and Decompositions of Games: Harmonic and Potential Games,” <i>Math. of OR</i> 36(3):474–503 (found 2026-05-21)	<i>formal antecedent</i> — Hodge-style decomposition of finite normal-form games into three orthogonal components; gives a coordinate-free organizing structure for the $R1$ -vs- $R2$ boundary cleaner than tier-wise enumeration. AAT adopts the decomposition as the canonical structural reading of the $R1/R2$ partition; the four-tier ladder is the AAT-internal consolidation across the decomposition plus the contraction-regime ($R0$) extension plus the mean-field-regime ($R3$) limit
Differentiable game dynamics with potential / Hamiltonian decomposition	Letcher, Balduzzi, Racanière, Martens, Foerster, Tuyls & Graepel 2019, “Differentiable Game Mechanics,” <i>JMLR</i> 20:1–40 (found 2026-05-21)	<i>formal antecedent</i> — lifts the Hodge-style decomposition from games to <i>dynamics on games</i> (potential component $\rightarrow$ convergent gradient flow; Hamiltonian component $\rightarrow$ cyclic flow). The $R1/R2$ partition is exactly the potential-vs-Hamiltonian dynamics distinction at the joint-best-response field level

<i>ASF concern</i>	<i>Prior-art language</i>	<i>Relationship / Positioning</i>
Cyclic-dynamics-as-natural-object in adversarial learning	Mertikopoulos, Papadimitriou & Piliouras 2017, “Cycles in Adversarial Regularized Learning,” <i>SODA 2018</i> (proceedings; arXiv:1709.02738, found 2026-05-22)	<i>formal antecedent (direct R2 predecessor)</i> — direct historical predecessor to the Piliouras-lab MCC line and to Cheung-Piliouras-Tao 2021: regret-minimization dynamics in adversarial regularized games exhibit <i>cycles</i> rather than convergence, and trajectory rate $\neq$ regret rate (a caveat load-bearing for #disc-identifiability-floor Instance 3 escape (e)’s “trajectory-rate not regret-rate” specification). Surface-back from the Working Notes BG2 caveat to the main Related Work table.
Population-games taxonomy (potential / stable / cyclic)	Sandholm 2010, <i>Population Games and Evolutionary Dynamics</i> , MIT Press (textbook; found 2026-05-21)	<i>formal antecedent</i> — population-level classification of strategic interaction by dynamic regime. The R0/R1/R2 partition tracks Sandholm’s “stable” / “potential” / “non-stable cyclic” taxonomy at the finite- N scale; the R3 tier is exactly Sandholm’s population-distribution limit
Evolutionary phase-portrait classification	Hofbauer & Sigmund 1998, <i>Evolutionary Games and Population Dynamics</i> , Cambridge University Press (textbook; found 2026-05-21)	<i>formal antecedent</i> — phase-portrait analysis of replicator dynamics distinguishing fixed-point convergence (R0/R1) from heteroclinic/cyclic structure (R2). The AAT regime axis adopts the phase-portrait reading as the structural shape of the R0/R1/R2 distinction
Cyclic dynamics as solution concept	Omidshafiei, Papadimitriou, Piliouras, Tuyls, Rowland, Lepiauw, Czarnecki, Lanctot, Perolat & Munos 2019, “ α -Rank: Multi-Agent Evaluation by Evolution,” <i>Scientific Reports</i> 9:9937 (found 2026-05-21)	<i>formal antecedent</i> — Markov-Conley-chain (MCC) solution concept treats cyclic dynamics as the natural object when Nash equilibria are absent or non-attracting. AAT adopts MCC as a sub-tier within R2’s distributional macro-state
Chain-recurrent set as outcome of non-converging dynamics	Papadimitriou & Piliouras 2018, “From Nash Equilibria to Chain Recurrent Sets,” <i>ACM Trans. on Economics and Computation</i> (found 2026-05-21 — local PDF ref/papadimitriou-piliouras-2018-nash-chain-recurrent.pdf)	<i>formal antecedent</i> — chain-recurrent set as the natural macro-state when fixed-point convergence fails. Provides the formal substrate for R2’s distributional macro-state on the dynamics side, complementing Hart-Mas-Colell’s CCE on the no-regret-learning side
Fundamental Theorem of Dynamical Systems (the universality behind chain-recurrence)	Conley 1978, <i>Isolated Invariant Sets and the Morse Index</i> , CBMS Regional Conf. Ser. Math. 38, AMS, §8.1 (found 2026-05-22)	<i>formal antecedent (direct citation)</i> — the universal decomposition of any flow on a compact invariant set into a chain-recurrent subflow plus a strongly-gradient-like quotient. Papadimitriou-Piliouras 2018’s chain-recurrent-set formalism builds on Conley’s theorem; the direct citation acknowledges Conley as the upstream mathematical source and as the universality backing R2’s macro-state classification. Also the cross-layer anchor for #result-certificate-existence’s R0-loss rung mixed-case (where the same universality localizes at the single-agent linearization).

ASF concern	Prior-art language	Relationship / Positioning
Coarse correlated equilibrium convergence	Hart & Mas-Colell 2000, “A Simple Adaptive Procedure Leading to Correlated Equilibrium,” <i>Econometrica</i> 68:1127–1150 (found pre-2026; cited in #deriv-strategic-composition)	<i>formal antecedent</i> — supplies the doubling-trick CCE convergence rate $O(1/\sqrt{T})$ that defines the R2 tier’s Lyapunov-equivalent guarantee
Mean-field games (MFG)	Lasry & Lions 2007, <i>Japanese J. Math.</i> 2:229–260; Huang, Malhamé & Caines 2006, <i>Comm. Info. Sys.</i> 6:221–252 (found 2026-05-21 by reference)	<i>formal antecedent</i> — supplies the R3 tier’s measure-space machinery. AAT does not currently have R3 machinery; the tier is named as the population-scope frontier per #deriv-strategic-composition Working Notes mean-field extension
Mechanism-design impossibility (designer-side regime-ascent floor)	Gibbard 1973, <i>Econometrica</i> 41:587–601; Satterthwaite 1975, <i>J. Econ. Theory</i> 10:187–217; Myerson & Satterthwaite 1983, <i>J. Econ. Theory</i> 29:265–281; Arrow 1951, <i>Social Choice and Individual Values</i> (found pre-2026; flagged in #deriv-strategic-composition Discussion)	<i>formal antecedent (designer-side floors)</i> — floor on what external designers can achieve in the $R2 \rightarrow R1$ or $R1 \rightarrow R0$ regime-ascent direction. Forward-references the queued #disc-implementation-impossibility meta-segment for the formal treatment

Search Log:

- 2026-05-21 (*targeted*, BG1 Pillar prior-art search): Verdict on T1 (four-tier classification as unifying construct): *Substantial Overlap with Named Prior Art* (High confidence). Verdict on T2 (macro-state-type / regime identity §C.1): *Substantial Novelty with Adjacent Literature* (Medium-High). Verdict on T3 (transition asymmetry §C.3): *Substantial Novelty with Adjacent Literature* (Medium). Verdict on T4 (identifiability floor §H.5): *Substantial Novelty with Adjacent Literature* (Medium). Most consequential finding: Candogan-Menache-Ozdaglar-Parrilo 2011 Hodge decomposition as the recognized prior art for the R1/R2 boundary (now in the Related Work table).
- 2026-05-21 (*planned*): Second-pass Undermind search for the broadened-Instance-3 + escape (e) to refine the R1-vs-R2 classifier-under-rate-class-observation claim if prior art surfaces.

Working Notes.

- **Sub-tier unification candidate.** Each regime has its own internal sub-classification (R0 Contraction ladder Tier 1/2/3; R1 potential/monotone/general; R2 CCE-convergence-rate; R3 MFG-uniqueness). Whether these unify across regimes under a single meta-axis — *Lyapunov-certificate-strength* (full-state Lyapunov / local-on-deviation Lyapunov / distributional-on-empirical Lyapunov / no Lyapunov) — is open. Candidate for a follow-on spike if the cross-regime meta-classification turns out to compose cleanly with existing tier machinery.
- **§C.3 cost-asymmetry derivation (was J9) + standalone-paper trajectory (was J7) — migrated 2026-05-25 to proper homes.** The §C.3 transition cost-asymmetry derivation with a structural-distance metric on game-rewiring space is registered at spikes/PROPOSED.md Tier 2 (added 2026-05-25; segment-rooted with detail home here in this Working Notes section). Landing the derivation would upgrade §”Transition asymmetry: descent is the default” from argumentation-grade to derived; estimated 2 cycles. The standalone-paper trajectory it was the prerequisite for is now at ~/src/ops/papers/deferred/dynamic-regime-axis.md (deferred until the derivation lands or is declined; BG1’s T1 Substantial Overlap verdict means the four-tier classification itself is not the publishable contribution, and the AAT-specific identities — §C.1 macro-state-type + §C.3 cost-asymmetry + §H.5 identifiability floor — are; without the §C.3 derivation, only §C.1 and §H.5 are substantive standalone contributions, and the paper is not yet competitive). Provenance trail: spikes/strategic-composition-class-3-attempt-2026-05-21/07-SOLIDIFIED-PLAN.md §S.2 originally framed these as the J7/J9 Joseph-reserved tail of the strategic-composition

cluster; with that cluster's primary work closed 2026-05-25 (Tracks A/B/C/D/CR/E all DONE), the two items are migrated to their proper navigators.

- **#disc-modularity-state-dynamics parallel-pattern.** The parallel-three-operation observation in §"Cross-axis interactions" connects to #disc-modularity-state-dynamics (landed 2026-05-24 closing the modularity-as-contested-property cycle per msc/modularity-cycle-plan-2026-05-09.md). With both meta-segments now stable, a higher-order meta-segment covering both Axes A and B at the parallel-pattern level becomes a candidate follow-on cycle if the cross-pattern proves load-bearing for downstream framework work; the current arrangement holds the parallel observation in each of the two without forcing the higher-order recognition.
- **A4 second-pass Undermind on broadened Instance 3 (resolved 2026-05-21).** Completed per 07-SOLIDIFIED-PLAN.md A4; report at ref/Unidentifiability_and_rate_class_prior_art.md^d. Verdict on both targets: *no exact prior-art hit; rich near-miss landscape*. The R1-vs-R2 classifier strengthened-consequence qualifier was softened from "unique broadly-available" to "primary AAT-framework" to reflect that adjacent passive diagnostics for *related* regime distinctions exist (potentialness — Bichler et al. 2025^b; convergence-vs-recurrence — Legacci et al. 2024^c; best-reply structure — Pangallo et al. 2017^d; algorithm-specific rate dichotomies — Anagnostides et al. 2022^e, Shi-Zhang 2019^f), but none names rate-class-as-passive-classifier-of-regime in the form escape (e) instantiates. The Related Work table in #deriv-regime-marginal-indistinguishability carries the full neighbor citations; #disc-identifiability-floor Instance 3 escape (e) and strengthened consequence carry the softened framing with neighbor-citation acknowledgment. Two important caveats from the BG2 verdict landed: (a) Mertikopoulos-Papadimitriou-Piliouras 2017 *Cycles in Adversarial Regularized Learning*^g shows regret-rate alone is too blunt — escape (e) must specify *trajectory* rate, not regret rate (now landed in Instance 3 escape (e) wording); (b) Bielecki et al. 2011^b is the closest structural analogue for the cross-topology marginal-matching witness at the multivariate Markov-chain level — the witness math in #deriv-regime-marginal-indistinguishability extends the established marginal-non-determining shape to the game-topology layer.
- **Primary-source verification spike queued (2026-05-21).** All BG2-introduced citations in this Working Note (marked with footnote [^hbg2-2026-05-21]) and in the cross-referenced #disc-identifiability-floor Instance 3 escape (e) and #deriv-regime-marginal-indistinguishability Related Work need primary-source verification before any further reliance. The characterizations attributed to those papers are currently synthesized from ref/Unidentifiability_and_rate_class_prior_art.md, not from direct reading. Verification targets for this segment specifically: Mertikopoulos-Papadimitriou-Piliouras 2017 (the trajectory-vs-regret-rate caveat — load-bearing for the softening of "unique" to "primary AAT-framework"); Bichler et al. 2025 (named-passive-diagnostic counterexample to "unique" qualifier); Legacci-Mertikopoulos-Pradelski 2024 (closest convergence-vs-recurrence taxonomy neighbor); Bielecki et al. 2011 (closest structural analogue, load-bearing for #deriv-regime-marginal-indistinguishability's positioning). Spike to be scheduled per Joseph's go.
- **Mechanism reduction question.** Whether the regime-axis transition asymmetry reduces to a special case of the modularity-state-dynamics three-operation pattern (Axis A) under a structural translation, or whether the two patterns are genuinely independent at the meta-architectural level, is open. The current §"Cross-axis interactions" reading treats them as parallel-but-independent; a unification attempt would be a separate spike.
- **R3 machinery.** AAT does not currently carry mean-field-equilibrium-regime machinery. R3 is named as the population-scope frontier; per #deriv-strategic-composition Working Notes, the mean-field-extension awaits Part III population-dynamics machinery (latent structural diversity; endogenous coupling; composition-transition dynamics). When R3 lands properly, this segment will be revisited to verify the four-tier ladder remains the right granularity (vs collapsing R2 + R3 into a unified distributional tier, or splitting R3 further by MFG sub-class).

^aThe paper and the characterization of its claims are inherited from a second-pass Undermind synthesis (search conducted 2026-05-21); the primary source has *not* been verified at landing time. A primary-source verification spike is queued — see Working Notes above.

^bThe paper and the characterization of its claims are inherited from a second-pass Undermind synthesis (search conducted 2026-05-21); the primary source has *not* been verified at landing time. A primary-source verification spike is queued — see Working Notes above.

^cThe paper and the characterization of its claims are inherited from a second-pass Undermind synthesis (search conducted 2026-05-21); the primary source has *not* been verified at landing time. A primary-source verification spike is queued — see Working Notes above.

^dThe paper and the characterization of its claims are inherited from a second-pass Undermind synthesis (search conducted 2026-05-21); the primary source has *not* been verified at landing time. A primary-source verification spike is queued — see Working Notes above.

^eThe paper and the characterization of its claims are inherited from a second-pass Undermind synthesis (search conducted 2026-05-21); the primary source has *not* been verified at landing time. A primary-source verification spike is queued — see Working Notes above.

^fThe paper and the characterization of its claims are inherited from a second-pass Undermind synthesis (search conducted 2026-05-21); the primary source has *not* been verified at landing time. A primary-source verification spike is queued — see Working Notes above.

^gThe paper and the characterization of its claims are inherited from a second-pass Undermind synthesis (search conducted 2026-05-21); the primary source has *not* been verified at landing time. A primary-source verification spike is queued — see Working Notes above.

^hThe paper and the characterization of its claims are inherited from a second-pass Undermind synthesis (search conducted 2026-05-21); the primary source has *not* been verified at landing time. A primary-source verification spike is queued — see Working Notes above.

Appendix AG

Formulation: Resource Budget

A depletable scalar reservoir whose drain rate rises with mismatch and whose level gates correction capacity — the minimal structure that makes “a degrading model is more expensive to run” a dynamical statement rather than an informal one.

The construction is minimal: a scalar resource budget $\mathcal{B}_t$ evolving under drain-versus-replenishment, with two introduced structural posits — (**A-cost**) the drain rate $c(\|\delta\|)$ is non-decreasing in mismatch (degrading models cost more to run), and (**A-gate**) the sector parameter is resource-gated via $\alpha(\mathcal{B}) = \alpha^{\max}\psi(\mathcal{B})$ with $\psi(0) = 0$ (the depleted agent loses corrective capacity entirely). Hard-budget regime has zero replenishment ($r_{\mathcal{B}} = 0$, a finite pool that only depletes); regenerative regime has $r_{\mathcal{B}} > 0$.

This segment is *exploratory off-spine*: no Part I/II/III result depends on it; it opens a resource-structure axis 6.7 explicitly records as an open scope item. The formulation is *conditional by nature* — (A-cost) and (A-gate) are modeling assertions about an agent’s physical realization, not derivations from core machinery, so the formulation’s epistemic ceiling stays at conditional regardless of downstream development. The companion derivation AH uses this formulation to close 15.2’s Effects-Spiral by an alternative mechanism (agent’s own correction-rate decay rather than adversary coupling growth).

Formal Expression.

AAT’s core machinery is resource-blind: the policy $\pi(M_t, G_t)$ acts at no modeled cost, and correction capacity (tempo $\mathcal{T}$, 3.8) does not deplete. 6.7 records this as an explicit open scope item — for resource-constrained agents (embodied controllers under torque/battery/episode-length limits; teams under headcount; any agent whose every corrective action spends an exhaustible pool) the formalism carries no state for the pool and no coupling from model quality to its drain. This formulation introduces the minimal such structure.

A scalar **resource budget** $\mathcal{B}_t \geq 0$ (calligraphic-scalar by the NOTATION convention exception, as for $\mathcal{T}$; distinct from the agent label B of 15.2 and from the strategy-edge set E). Its evolution:

Definition (resource-budget).

$$\frac{d\mathcal{B}}{dt} = -c(\|\delta\|) + r_{\mathcal{B}} \quad (\text{AG.1})$$

where δ is the agent’s mismatch (3.4), $c(\cdot) \geq 0$ is the **correction-cost rate**, and $r_{\mathcal{B}} \geq 0$ is the **replenishment rate**. The hard-budget regime is $r_{\mathcal{B}} = 0$ (a finite pool that only depletes — a combat episode’s battery, a fixed torque-integral, a bounded step count); the regenerative regime is $r_{\mathcal{B}} > 0$.

The formulation is fixed by two structural posits — introduced, not derived from existing AAT.

c is non-decreasing in $\|\delta\|$, with $c(0) > 0$ (running the loop at all costs something) and c strictly increasing where $\|\delta\| > 0$. Minimal concrete form:

Assumption (A-cost: cost rises with mismatch).

$$c(\|\delta\|) = c_0(1 + \beta_{\mathcal{B}} \|\delta\|), \quad c_0 > 0, \beta_{\mathcal{B}} \geq 0. \quad (\text{AG.2})$$

Motivation, not derivation: a model that is wrong by δ actuates partly in wrong directions — the agent pays for the corrective action *and* for the wasted component, and re-observes more often to recover. Degradation is literally more expensive to carry. $\beta_{\mathcal{B}} = 0$ recovers the resource-blind special case (cost independent of model quality).

The sector/correction rate the persistence machinery uses — α in [G](#), equal to $\mathcal{T}$ in the canonical epistemic instantiation ([3.8](#), [4.5](#)) — is throttled by available resource:

Assumption (A-gate: tempo is resource-gated).

$$\alpha(\mathcal{B}) = \alpha^{\max} \psi(\mathcal{B}), \quad \psi : \mathbb{R}_{\geq 0} \rightarrow [0, 1] \text{ non-decreasing, } \psi(0) = 0, \psi(\mathcal{B}) \rightarrow 1 \text{ as } \mathcal{B} \rightarrow \infty. \quad (\text{AG.3})$$

Motivation, not derivation: with the pool exhausted the agent cannot run its observe–update–actuate loop at full rate (fewer sensor sweeps, slower control cycle, fewer remaining episode steps). $\psi \equiv 1$ recovers the resource-blind special case (constant α , exactly today’s template).

Together: $\mathcal{B}_t$ is the new state; (A-cost) couples model quality *into* its drain; (A-gate) couples its level *into* correction capacity. These two couplings are exactly what close the otherwise-open feedback in [15.2](#)’s Effects-Spiral corollary — the consequence is derived in [AH](#).

Epistemic Status.

Conditional. This is a representational choice (type: formulation) — several resource models would fit (per-action cost vectors, capacity constraints on the action set, queueing-theoretic service limits); $\mathcal{B}_t$ with (A-cost)/(A-gate) is the *minimal* one that makes the model-quality $\rightarrow$ drain $\rightarrow$ capacity loop dynamical with one added scalar and two monotone couplings. The defined objects ($\mathcal{B}_t$, the depletion law) are well-posed; the *content* is the two posits, which are introduced premises, honest as modeling commitments and not consequences of any prior segment. Anything downstream is conditional on them.

Max attainable: conditional. This ceiling is intrinsic: (A-cost) and (A-gate) are claims about a particular agent’s physical realization, not about the formalism, so no amount of derivation promotes the formulation itself beyond conditional. A *specific* agent class can satisfy them by construction (e.g. a controller whose actuation energy is measured and whose control-loop frequency is power-throttled), at which point downstream results inherit that class’s tighter status — but the formulation as stated stays conditional by nature, the same way [6.7](#)’s other scope items do.

This segment is an exploratory branch: it opens a resource-structure axis AAT has deliberately lacked. No spine segment depends on it; it is not a prerequisite of any Part I/II/III result.

Discussion.

What the two posits buy, minimally. Everything load-bearing downstream uses only *monotonicity*: (A-cost) that drain is non-decreasing in mismatch, (A-gate) that capacity is non-decreasing in budget with $\psi(0) = 0$. The explicit linear c and the shape of ψ are illustrative; the qualitative consequences in AH survive any (c, ψ) obeying the monotonicity and the endpoint $\psi(0) = 0$. The endpoint matters: $\psi(0) = 0$ encodes “an empty pool cannot correct at all,” which is what makes exhaustion *terminal* rather than merely *slowing*.

Why a new scalar rather than folding into ρ or $\mathcal{T}$. Resource is not disturbance: ρ is mismatch injected by the world, independent of how well the agent is doing; $\mathcal{B}_t$ drains *because* the agent is doing badly (A-cost). Nor is it tempo: $\mathcal{T}$ is a rate the agent has; $\mathcal{B}_t$ is a stock it spends to exercise that rate. Collapsing resource into either erases the very coupling that makes the Effects-Spiral close — a degrading model raising its own future cost is a stock-flow statement, not a rate statement. The new scalar is the minimum that keeps stock and flow distinct.

Relationship to the adaptive reserve $\Delta\rho^*$. G’s adaptive reserve $\Delta\rho^* = \alpha R - \rho$ is a *margin* (how much extra disturbance the agent can absorb at its current rate); $\mathcal{B}_t$ is a *fuel* (how much correction it can still afford to perform). They are independent: an agent can have comfortable margin and an almost-empty pool, or ample fuel and zero margin. The resource-bounded result is precisely about a system that is margin-safe at full fuel and still dies as the fuel drains the margin away — see AH.

Domain instances. Embodied combat/contested control under a battery or torque-integral budget (the case that motivated this branch); a development team under a fixed person-hour pool where a worsening codebase model raises rework cost; an incident-response org whose escalating confusion burns the on-call budget faster. In each, the resource-blind special case ($\beta_{\mathcal{B}} = 0, \psi \equiv 1$) is the agent for whom the pool is effectively infinite relative to the engagement — and there the existing constant- α machinery is exactly right, which is the honest boundary of this branch’s relevance.

Working Notes.

- The regenerative regime $r_{\mathcal{B}} > 0$ with $\mathcal{B}_t$ bounded above (a finite-capacity rechargeable pool) is the genuinely open dynamical case — it is a time-varying- α template instantiation in full and may be quasi-stationary-distribution / absorbing-barrier-flavored (cf. the open second no-go in the continuity-persistence reasoning trail). AH treats only $r_{\mathcal{B}} = 0$ rigorously; the regenerative case is flagged there and not attempted.
 - Reasoning trail: spikes/fight/03-energy-bound-effects-spiral.md (the Φ -prompted derivation attempt that produced this branch) and spikes/fight/99-verdict.md §3.
 - Open: whether $\mathcal{B}_t$ should be vector-valued (distinct pools — energy vs compute vs attention — with distinct c, ψ). Scalar is the minimal first cut; the persistence machinery it feeds is itself scalar-default with a tensor extension (3.8), so a vector $\mathcal{B}$ would compose with that tensor form. Not pursued here.
 - A NOTATION.md entry for $\mathcal{B}_t$ is added under a resource section; per the NOTATION drift caveat the segment, not the index, is ground truth.
-

Appendix AH

Derived: Resource-Bounded Destabilization

A hard-budget agent self-depletes to certain finite-time destabilization against even a *constant*-effectiveness adversary — closing 15.2’s Effects-Spiral not by formalizing its open coupling term but by making that term unnecessary, via the decaying- α instantiation of **G**.

Formal Expression.

15.2’s Effects-Spiral corollary ($\|\delta_B\| \uparrow \Rightarrow \text{erratic action} \Rightarrow \gamma_A \uparrow \Rightarrow \rho_B \uparrow \Rightarrow \|\delta_B\| \uparrow$) is *discussion-grade* there because the link “degrading model $\Rightarrow$ stronger adversary coupling” requires specifying how model degradation feeds back into the dynamics, which that segment’s machinery does not carry. The resource-budget formulation (**AG**) supplies that feedback through a different and cleaner channel — the agent’s *own* correction rate, not the adversary’s coupling.

Setup: the resource-coupled adversarial instantiation. Instantiate **G** with the target agent *B*’s mismatch as state variable, $\xi = \delta_B$, exactly as 15.2 does, but with the sector parameter made resource-gated per **AG** (A-gate):

$$\frac{d\delta_B}{dt} = -\alpha_B(\mathcal{B}) \delta_B + w_B(t), \quad \|w_B(t)\| \leq \rho_B^{\text{eff}} = \rho_{B,\text{base}} + \gamma_A \mathcal{T}_A, \quad \frac{d\mathcal{B}}{dt} = -c(\|\delta_B\|) + r_B, \quad \alpha_B(\mathcal{B}) = \alpha_B^{\text{max}} \psi(\mathcal{B}),$$

with ρ_B^{eff} the coupling-amplified disturbance of 15.2 (treated there, and here, with $\gamma_A, \mathcal{T}_A$ exogenous) and c, ψ as posited in **AG** (A-cost), (A-gate). The novelty is the *coupled pair*: δ_B and $\mathcal{B}$ co-evolve, with model quality draining the budget and the budget gating correction.

Regime split.

Budget-sufficient regime. If $\mathcal{B}_t$ stays large enough that $\psi(\mathcal{B}_t) \approx 1$ over the engagement (the pool is effectively infinite relative to the integrated cost), then $\alpha_B(\mathcal{B}_t) \approx \alpha_B^{\text{max}}$ is constant and the system is 15.2’s constant- α instantiation, unchanged. The resource structure adds nothing here. This is the honest boundary: the resource-blind special case ($\psi \equiv 1$) is exactly today’s machinery, and it is correct there.

Derived (resource-regime-split, from form-resource-budget + sector-persistence-template).

Budget-scarce hard regime ($r_B = 0$). With no replenishment, $d\mathcal{B}/dt = -c(\|\delta_B\|) \leq -c(0) < 0$ whenever correction runs, so $\mathcal{B}_t$ is strictly decreasing and reaches any level in finite time. Then $\alpha_B(\mathcal{B}_t)$ is a **time-varying, monotonically-decaying** sector parameter — exactly the structure **G**’s Epistemic Status flags as requiring additional machinery, and the structure 9.10 already instantiates (there $\alpha_\Sigma = 1/(n+1)$ decays with experience; persistence then requires the decay be counteracted faster than ρ_Σ/R_Σ). The resource-bounded case is the *same template slot* with budget-depletion in the role experience-accumulation plays for strategy persistence.

The result.

Define the **critical budget** $\mathcal{B}_{\text{crit}}$ by the persistence boundary of G (Model D) at the coupling-amplified disturbance:

Derived (Conditional on A-cost, A-gate; hard regime $r_B = 0$).

$$\alpha_B(\mathcal{B}_{\text{crit}}) = \frac{\rho_B^{\text{eff}}}{R_B} \iff \psi(\mathcal{B}_{\text{crit}}) = \frac{\rho_B^{\text{eff}}}{\alpha_B^{\text{max}} R_B}, \quad (\text{AH.1})$$

and the destabilization **hitting time** $\tau = \inf\{t : \mathcal{B}_t = \mathcal{B}_{\text{crit}}\}$. Then, in the hard regime with correction ongoing:

1. **Certain finite-time destabilization.** $\mathcal{B}_t \downarrow$ strictly, so $\tau < \infty$ for any finite $\mathcal{B}_0$. Because $\psi(0) = 0$, once $\mathcal{B}_t < \mathcal{B}_{\text{crit}}$ the persistence condition $\alpha_B(\mathcal{B}_t) > \rho_B^{\text{eff}}/R_B$ fails and 15.2's destabilization threshold is crossed. **The static persistence inequality holding at $t = 0$ does not prevent this** — an agent that would persist forever at full budget ($\alpha_B^{\text{max}} R_B > \rho_B^{\text{eff}}$) still destabilizes once the fuel drains the rate below threshold.
2. **The spiral, derived.** Substituting (A-cost), a mismatch excursion accelerates its own destabilization: $\|\delta_B\| \uparrow \Rightarrow c(\|\delta_B\|) \uparrow \Rightarrow \dot{\mathcal{B}}$ more negative $\Rightarrow \alpha_B(\mathcal{B}) \downarrow \Rightarrow (\alpha_B R_B - \rho_B^{\text{eff}}) \downarrow \Rightarrow \|\delta_B\| \uparrow$, which brings τ forward: a transient noise burst in w_B permanently advances the destabilization time. This is 15.2's Effects-Spiral, now a derived consequence of (A-cost)/(A-gate) rather than a discussion-grade schematic.
3. **The adversary's coupling need not grow.** The spiral closes through α_B *decaying* with γ_A held **constant**. The Effects-Spiral's substantive content — collapse invisible to the static threshold — is captured *without* the unspecified $\gamma_A(\|\delta_B\|)$ leg 15.2 could not formalize. Adversarial survival reduces to a **budget-versus-engagement-length race**: B persists iff the engagement ends before τ .

What Is Derived vs. What Is Chosen.

Table AH.1

<i>Property</i>	<i>Source</i>	<i>Strength</i>
Budget-sufficient regime $\equiv$ constant- α template, unchanged	G (status exact), $\psi \equiv 1$ limit	Inherited exact
Resource-scarce case is the decaying- α template slot (the reduction)	G time-varying- α + 9.10 precedent	Derived, conditional on (A-cost),(A-gate)
Certain finite-time destabilization, $r_B = 0$	Monotone-budget + Lyapunov persistence-boundary crossing	Derived (conditional)
Effects-Spiral as derived feedback; τ advanced by excursions	(A-cost) substituted into the depletion law	Derived (conditional)
Adversary coupling need not grow (γ_A constant suffices)	Corollary of the above	Derived (conditional)
Orthogonal to — not a resolution of — the deferred symmetric joint-Jacobian problem	Comparison of which term carries the feedback (α_B vs γ_A)	Verified relationship
(A-cost), (A-gate)	AG	Introduced posits, not derived
Regenerative regime $r_B > 0$	—	Open (not attempted)

Epistemic Status.

Conditional on AG's two introduced posits (A-cost),(A-gate) and the hard regime $r_B = 0$. Given those, the chain is tight: the reduction to the decaying- α slot of G is exact-flavored (it is that template's own machinery, the same instantiation 9.10 makes for experience-decay), and certain finite-time destabilization follows by a monotone-budget + Lyapunov-boundary argument with no further assumption. The status is conditional rather than exact because the load-bearing premises are modeling posits about an agent's physical realization, not consequences of prior segments — inherited honestly from AG's own ceiling.

Maxattainable: conditional for the formulation-dependent result; a specific agent class that satisfies (A-cost)/(A-gate) by construction (measured actuation energy, power-throttled control loop) lifts *its* instance to that class's tighter status, but the general statement stays conditional by the nature of its premises.

Scope honesty — what this does not do. It does *not* resolve the symmetric joint-best-response / joint-Jacobian equilibrium problem 15.2 defers to 16.1. That deferred problem is about both agents' mismatch co-evolving and specifically about the adversary's $\gamma_A(\|\delta\|)$ leg; this result closes a *different* leg (the target's own resource-gated α_B) and is single-agent-with-resource. The two are orthogonal: this result needs no claim about γ_A growing, and makes none. Reading it as having cracked the symmetric equilibrium problem would be an overclaim — it routes *around* that problem by adding resource structure, at the cost of two posits AAT otherwise lacks.

The regenerative regime ($r_B > 0$) is genuinely open and not attempted here (see Working Notes).

Discussion.

Why this is a strengthening, not a softening. The honest move when facing 15.2's discussion-grade Effects-Spiral is not “leave it discussion-grade.” It is to ask whether the spiral can be made a theorem. It can — in the resource-bounded regime, by a reduction to existing exact machinery — and the resulting statement is *sharper* than the original conjecture: it eliminates the need for the unformalized $\gamma_A(\|\delta_B\|)$ term rather than supplying a guessed functional form for it. A constant-effectiveness adversary against a hard-budget target is already sufficient for certain finite-time collapse. The strengthen-before-soften discipline produced a cleaner result than the discussion-grade original gestured at.

The reframing is the payoff. “Can the adversary destabilize me?” becomes “does the engagement outlast my fuel?” Persistence under a competent constant adversary is no longer a static inequality on (α, ρ, R) but a race between the hitting time τ and the engagement horizon — and τ is endogenous, advanced by every mismatch excursion the agent suffers. This is the precise sense in which a degrading embodied agent under hard physical limits “spirals”: not because the opponent gets stronger, but because being wrong is expensive and the agent is spending a finite pool to be wrong in.

Relationship to 9.10. That schema and this result are the two non-epistemic instantiations of G's *time-varying- α* extension. Strategy persistence's α_Σ decays with accumulated experience and is rescued by forgetting; here α_B decays with spent budget and is rescued only by replenishment or by ending the engagement before τ . Same template slot, different decay driver, structurally parallel rescue conditions. The parallel is itself evidence the template's time-varying- α extension is the right general object, with these as two of its instances.

Findings.

Resource Depletion Closes the Effects Spiral by Eliminating Its Open Term. Brief: A robot fighting another robot on a finite battery does not lose because its opponent gets cleverer as it falters — it loses because faltering is expensive. A wrong model wastes actuation, the wasted actuation drains the battery, the drained battery slows the control loop, the slower loop lets the model fall further behind: a self-reinforcing collapse that needs no help from the adversary. The everyday version: a tired driver makes more mistakes, each mistake costs more energy to recover from, and the exhaustion that follows makes the next mistake likelier — the spiral is in the fuel, not in the road. AAT had this spiral only as an informal sketch because it could not say how a worse model makes an opponent stronger. The resource

view removes the question: the opponent's strength can be held fixed; a finite, mismatch-sensitive fuel pool is enough to force certain collapse in finite time, with the only open question being whether the fight ends before the fuel does.

Impact: Converts 15.2's Effects-Spiral corollary from discussion-grade to conditional-derived — not by formalizing its unspecified adversary-coupling term but by making that term unnecessary, via a reduction to the decaying- α slot of G (the same machinery 9.10 uses for experience-decay). Reframes adversarial survival under a competent constant adversary from a static (α, ρ, R) inequality to a budget-versus-engagement-length race with an endogenous, excursion-advanced hitting time. Opens the resource-structure axis 6.7 records as an explicit gap, as an exploratory branch off the spine. Explicitly does *not* resolve the deferred symmetric joint-Jacobian problem (16.1) — it is orthogonal, and the Finding marks that boundary so the result is not over-read.

Novelty Claim: *Claim differentiation* on 15.2's Effects-Spiral: the contribution is showing the spiral becomes a derived finite-time-destabilization result once a minimal resource state is added, and that the strengthening *removes* rather than supplies the previously-open coupling functional form. The sector-Lyapunov and decaying- α machinery are AAT-internal and standard; resource-constrained control and energy-bounded games are established external territory; the differentiation is the specific recognition that the energy bound is the feedback variable that closes AAT's own named-open spiral by eliminating its open leg, conditional on two explicit posits.

Related Work:

Table AH.2

<i>ASF concern</i>	<i>Prior-art language</i>	<i>Relationship / Positioning</i>
Resource/energy budgets shaping control behavior	Energy-bounded / resource-constrained control and RL (broad external literature)	<i>adjacent literature</i> — supplies the modeling milieu (depletable budgets, cost-of-control); AAT's contribution is the coupling of model quality into drain and budget into the sector parameter, closing a specific AAT-internal spiral
Decaying-gain $\rightarrow$ eventual instability	9.10 (experience-decay forgetting prerequisite)	<i>formal antecedent (AAT-internal)</i> — the same time-varying- α template slot; resource-depletion is its second non-epistemic instantiation
Effects-Spiral as informal panic mechanism	15.2 Corollary (discussion-grade)	<i>formalized by this finding</i> — the prior segment's open feedback term is here made unnecessary, not filled in

Search Log: - 2026-05-19 (*intuition-only*): The resource-constrained-control and energy-bounded-games literatures are large and were not searched; the AAT-internal claim (energy bound as the feedback variable that closes the Effects-Spiral by eliminating its open $\gamma_A(\|\delta\|)$ leg, conditional on (A-cost)/(A-gate)) has not been searched against that literature. The constituent machinery (sector-Lyapunov, decaying- α) is AAT-internal and standard; the differentiation claim is plausibly distinctive but unverified. Targeted future search: cost-of-control / energy-aware MPC, budgeted-RL regret under adversaries, lazy-control and event-triggered control under depletable actuation. Pre-search expectation: constituent moves well-precedented externally; the AAT-internal “closes its own named-open spiral by removal” framing is the candidate-distinctive part.

Working Notes.

- Reasoning trail: spikes/fight/03-energy-bound-effects-spiral.md (the Φ -prompted derivation attempt that produced this branch) and spikes/fight/99-verdict.md §2–3. The math lands here per math-lives-in-segments; the spike is the trail only.
- Open — regenerative regime.** $r_B > 0$ with a finite-capacity pool is the full time-varying- α ODE and is not attempted. It is plausibly quasi-stationary-distribution / absorbing-barrier-flavored (structural cousin of the open second no-go in the continuity-persistence reasoning trail). The natural next spike: classify whether bounded replenishment can guarantee $\mathcal{B}_t >$

$\mathcal{B}_{\text{crit}}$ indefinitely or only postpone τ in expectation (a Feller-test / absorbed-chain boundary question, exactly parallel to the continuity-persistence barrier-reachability open problem).

- **Open — rate constants.** τ as an explicit function of $(\mathcal{B}_0, \beta_{\mathcal{B}}, \psi, \text{engagement length})$ is not computed; only its finiteness and monotone advancement under excursions are derived. A closed-form τ under the linear c and a specific ψ (e.g. saturating $\psi(\mathcal{B}) = \mathcal{B}/(\mathcal{B} + \mathcal{B}_{1/2})$) is a tractable follow-on if the branch is pursued.
 - Exploratory-branch status: no spine segment depends on this or on [AG](#); OUTLINE marks both exploratory. Promotion past draft is gated on Joseph electing to develop the resource axis (regenerative case + a worked agent class satisfying (A-cost)/(A-gate) by construction would be the path to lifting an instance above conditional).
-

Appendix AI

Result: Contraction Template

#result-sector-persistence-template states AAT's persistence arguments with a Euclidean sector condition (T2) matched to a quadratic Lyapunov in Euclidean norm. Generalizing the sector inequality to a **contraction-metric condition** (Lohmiller & Slotine 1998) preserves the template's ultimate-bound results while extending its coverage in three directions: (a) sub-scope α gains natural non-Euclidean metrics that remove condition-number penalties currently visible in 4.3; (b) two additional sub-scope β items (PID-bounded-plant; non-convex-within-basin) promote to derived sub-scope metric- α_2 under explicit conditions; (c) composition acquires **topology-indexed closure results** (parallel / hierarchical / small-gain-feedback) from Slotine 2003 that generalize #deriv-critical-mass-composition's matched-symmetric-Tier-1 closed form to heterogeneous sub-agents. This segment states the generalization once in parameter-free form so that each lifting citation can reference it and specify only what varies.

#result-sector-persistence-template remains as the $M = I$ Euclidean specialization.

Formal Expression.

Preconditions.

Let $\xi(t) \in \mathbb{R}^n$ evolve under $\dot{\xi} = -F(\xi, t) + w(t)$ where F is C^1 in ξ and continuous in t . Let $M : \mathbb{R}^n \times \mathbb{R}_+ \rightarrow \mathbb{S}_{++}^n$ be a smooth symmetric positive-definite matrix-valued function with uniform conditioning:

Template preconditions (contraction-template).

$$m_1 I \preceq M(\xi, t) \preceq m_2 I \quad \text{for all } \xi \in \mathcal{B}_R, t \geq 0 \quad (\text{M0})$$

with constants $0 < m_1 \leq m_2 < \infty$.

(CT1) **Zero correction at zero state.** $F(0, t) = 0$ for all t .

(CT2) **Local differential-contraction condition.** There exist $\lambda > 0$ and $R > 0$ such that for all $\xi \in \mathcal{B}_R, t \geq 0$:

$$\dot{M}(\xi, t) + M(\xi, t) \frac{\partial F}{\partial \xi}(\xi, t) + \left(\frac{\partial F}{\partial \xi}(\xi, t) \right)^T M(\xi, t) \geq 2\lambda M(\xi, t). \quad (\text{CT2})$$

(CT3) **Bounded disturbance.** Either Model D ($\|w(t)\| \leq \rho_\xi$) or Model S ($w(t)$ Wiener-process increment with $\mathbb{E}[\|w(t)\|^2] = \sigma_\xi^2$).

Ultimate bound — Model D.

Under the preconditions with $V(\xi, t) = \xi^T M(\xi, t) \xi$, the state is ultimately bounded:

Result (contraction-template-D), conditional on (M0), (CT1)–(CT3-D).

$$\limsup_{t \rightarrow \infty} \|\xi(t)\| \leq \frac{\rho_\xi}{\lambda} \sqrt{\frac{m_2}{m_1}}. \quad (\text{CT-D})$$

Structural persistence (the ultimate bound fits within the contraction region $\mathcal{B}_R$) requires

$$\lambda R \sqrt{m_1/m_2} > \rho_\xi. \quad (\text{CT-D-persist})$$

Proofsketch. Compute $\dot{V} = \xi^T \dot{M} \xi - 2\xi^T M F + 2\xi^T M w$. Integrate (CT2) along the ray $s\xi$, $s \in [0, 1]$, using $F(0) = 0$: $2\xi^T M F(\xi, t) \geq 2\lambda \xi^T M \xi - \xi^T \dot{M} \xi$. Substituting: $\dot{V} \leq -2\lambda V + 2\xi^T M w$. Cauchy-Schwarz in M -inner product + (M0) gives the affine-contraction bound on $W = \sqrt{V}$; conversion to Euclidean norm via (M0) yields (CT-D). $\square$

Ultimate bound — Model S (with Itô correction).

Under stochastic disturbance $d\xi = -F(\xi, t) dt + \sigma_\xi dW_t$ and a metric M satisfying the Itô-correction bound $\frac{1}{2}\sigma_\xi^2 \text{tr}(M + \text{Hessian correction in drift direction}) \leq c_{\text{Itô}} m_2$ locally (automatic for state-independent M ; bounded by M 's Hessian otherwise), the stopped process satisfies

Result (contraction-template-S), conditional on (M0), (CT1)–(CT3-S), Itô-compatible metric.

$$\mathbb{E}[V(\xi(t \wedge \tau_R), t \wedge \tau_R)] \leq V(\xi(0), 0) e^{-2\lambda t} + \frac{n\sigma_\xi^2 c_{\text{Itô}} m_2}{2\lambda}, \quad (\text{AI.1})$$

where $\tau_R = \inf\{t : \|\xi(t)\| > R\}$. Mean-square structural persistence requires

$$\lambda R^2 > \frac{n\sigma_\xi^2 c_{\text{Itô}} m_2^2}{2m_1}. \quad (\text{CT-S-persist})$$

This is the natural extension of #deriv-sector-condition's Prop A.1S region condition (Khasminskii 2012 ch. 5) to weighted metrics. When $M = I$ (Euclidean), $c_{\text{Itô}} = 1$, and the result reduces to Prop A.1S.

Recovery of #result-sector-persistence-template. When $M \equiv I$ (Euclidean metric, time-independent), (CT2) reduces to

$$-\frac{\partial F}{\partial \xi} - \left(\frac{\partial F}{\partial \xi}\right)^T \leq -2\lambda I, \quad (\text{AI.2})$$

which is the **incremental Euclidean sector condition** (strong monotonicity of F in Euclidean norm, equivalently **DA2'-inc** in #form-composition-closure's bridge-lemma language). This implies the one-point sector condition $\xi^T F(\xi) \geq \lambda \|\xi\|^2$ by integration along the ray $s\xi$, so $\alpha = \lambda$ when $M = I$. The template's ultimate bound ρ_ξ/α is reproduced with $m_1 = m_2 = 1$.

The structural observation: (CT2) at $M = I$ is strictly stronger than (T2) — it requires the Jacobian's symmetric part to be uniformly positive definite everywhere, matching DA2'-inc rather than (T2). The contraction formulation therefore **collapses (T2) and DA2'-inc into one condition** at the cost of requiring differential (Jacobian-level) rather than

integral (inner-product-at-a-point) information. This is one of the template’s structural contributions: the same condition that runs the template runs the bridge-lemma that #form-composition-closure depends on, so the gap between single-agent persistence and composite tracking faithfulness closes when (CT2) holds.

Sub-scope metric- α_1 / metric- α_2 / metric- β partition. The contraction formulation refines #deriv-sector-condition’s A2’ sub-scope partition into three tiers:

Sub-scope metric- α_1 (Euclidean metric, AAT-internally derived via DA2’-inc $\equiv$ (CT2) with $M=I$). Scalar Kalman, Euclidean strongly-convex gradient, L2-regularized convex, linear-PD-symmetric. These cases carry (CT2) with $M = I$ which is exactly DA2’-inc — a condition #form-composition-closure has required at the composite level all along. AAT’s existing commitment already forces (CT2)-Euclidean at these cases; the template surfaces it.

Sub-scope metric- α_2 (non-Euclidean metric, derived under explicit conditions). Five cases lift here:

- *Matrix Kalman under information metric* $M = (P^-)^{-1}$: contraction rate $\lambda = \lambda_{\min}(H^T R^{-1} H)/2$ on the observable subspace; Euclidean $\kappa(P^-)$ degradation removed. Under (PI)/Čencov (see #der-gain-sector-bridge “Fisher-metric cases under parameterization-invariance”), the information-metric form is *uniquely forced* — AAT-internally derived rather than theorem-imported.
- *Exponential-family natural parameters under Fisher metric* $M = I(\theta)$: contraction rate $\lambda = \eta$ globally on the interior of the natural-parameter domain (Fisher-conditioning degradation removed). Also AAT-internally forced under (PI)/Čencov.
- *Hessian-metric strongly-convex* $M = \nabla^2 L$: contraction rate $\lambda = \eta$ under bounded metric-derivative in the drift direction. Theorem-imported (Hessian metric chosen to match loss curvature; no AAT-internal axiom forces the specific coordinate).
- *Linear-Hurwitz-non-symmetric under Lyapunov metric*: M solves $MA + A^T M = -Q$ for $Q \succ 0$. Contraction rate $\lambda = \lambda_{\min}(Q)/(2\lambda_{\max}(M))$ in the M -metric. **New coverage:** asymmetric-stable linear corrections where Euclidean A2’ fails. Theorem-imported (Lyapunov equation construction is standard; no AAT-internal axiom forces Q).
- *PID with bounded plant nonlinearity under Lyapunov metric*: derived under plant-Lipschitz bound $L_p < \lambda_{\min}(Q)/(2\lambda_{\max}(M))$. **Promotion from sub-scope β .** Theorem-imported for the Lyapunov-metric construction; specializations include SPR-tuned PID (phase margin as sector constant; see #der-gain-sector-bridge Verified Instances).

Sub-scope metric- β (contraction-metric formulation fails). Four cases:

- *Variational / amortized / approximate posteriors*: contraction to the *projected* target (best-in-class q^*) is derivable, but the projection error $\text{KL}[q^* \| p_t]$ to the true posterior is a residual disturbance that contraction machinery cannot address.
- *Rule-based / symbolic / discontinuous updates*: non-smooth F ; (CT2) requires C^1 . Piecewise-smooth extensions (Di Bernardo et al. 2014) cover switched systems but not general rule-based reasoning.
- *Severely misspecified agents*: contraction to a wrong target is still wrong. Metric choice is silent on target validity.
- *Per-step SGD / human judgment*: noise-is-disturbance treatment identical to the Euclidean formulation; no improvement from metric choice.

This is **the seventh ladder** in #disc-separability-pattern (A2’-scope): metric- α_1 separable-core / metric- α_2 structured-repair / metric- β general-open.

Compositional contraction.

When sub-agents individually satisfy the preconditions with metrics M_i and rates λ_i , the composite satisfies (CT1)–(CT3) under specific topologies:

Result (compositional-contraction), conditional on Slotine 2003 Thm 1–3 + per-sub-agent (CT1)–(CT3).

(CC-parallel) Parallel composition. $\dot{x} = (\dot{x}_1, \dot{x}_2)^T$ with independent $(\dot{x}_i = -F_i(x_i) + w_i)$. The composite contracts under $M_c = \text{blockdiag}(M_1, M_2)$ at rate $\lambda_c = \min(\lambda_1, \lambda_2)$. Recovers #der-team-persistence’s weakest-link bound; cooperative coupling improves effective disturbance via the existing signed-coupling mechanism.

(CC-cascade) Hierarchical / lower-triangular cascade. If $\dot{x}_1 = f_1(x_1)$ contracts at λ_1 and $\dot{x}_2 = f_2(x_1, x_2)$ contracts in x_2 at λ_2 uniformly in x_1 with coupling gain $\|\partial f_2 / \partial x_1\| \leq k$, the cascade contracts. Rate bounded below by $\min(\lambda_1, \lambda_2)$ up to coupling-gain-dependent adjustment.

(CC-feedback) Negative feedback with bounded loop gain (Slotine 2003 Thm 3). Two systems with rates λ_1, λ_2 connected by negative feedback with loop gains k_{12}, k_{21} — the closed loop contracts iff

$$k_{12}k_{21} < 4\lambda_1\lambda_2. \quad (\text{CC-feedback})$$

Heterogeneous critical-mass — (CM2-M).

Specializing (CC-feedback) to the signed-coupling structure of #deriv-critical-mass-composition for heterogeneous Tier-1M sub-agents with metric-contraction rates λ_1, λ_2 , coordination costs C_1, C_2 , and feedback loop gains k_{12}, k_{21} :

Derived (CM2-M), from (CC-feedback) + ‘#deriv-critical-mass-composition’ signed-coupling structure.

$$(\lambda_1 - C_1)(\lambda_2 - C_2) > k_{12}k_{21}/4. \quad (\text{CM2-M})$$

Specializing further to the matched-symmetric case ($\lambda_1 = \lambda_2 = \lambda, C_1 = C_2 = C, k_{12} = k_{21} = k$): $(\lambda - C)^2 > k^2/4 \iff \lambda - C > k/2$, which matches #deriv-critical-mass-composition’s (CM2) with $k = 2\gamma\mathcal{T}/R$ (up to normalization). **Heterogeneous-architecture composites now have a closed-form critical-mass inequality** where the matched-symmetric case was the only closed form previously available; #deriv-critical-mass-composition’s §6.1 obstruction on heterogeneous composites is thereby closed for the Tier-1M case.

Epistemic Status.

Conditional. Conditional on (M0) + (CT1)–(CT3) and the specific metric M chosen for each sub-scope α entry. The (CT-D) and (CT-S) ultimate-bound results are standard Lohmiller-Slotine theory (Lohmiller & Slotine 1998; Slotine 2003) specialized to AAT’s per-instance state-variable choices; the proof technique is integration along the ray $s\xi$ followed by affine contraction. The contraction-metric primitive (Riemannian M) is standard in nonlinear control theory (Lohmiller-Slotine 1998; Forni & Sepulchre 2014 for the Finsler extension); AAT’s contribution is the AAT-internal mapping of M -choices to sub-scope α_1/α_2 and the recognition that the compositional theorems (CC-parallel / CC-cascade / CC-feedback / CM2-M) directly generalize Part III’s existing closure machinery.

The Jacobian-level observation. #form-composition-closure’s DA2’-inc condition is equivalent to (CT2) with $M = I$ for $C^1 F$ on convex domains. AAT has been carrying the Jacobian-level Euclidean contraction condition at the composite level all along under the name DA2’-inc; the contraction template’s Euclidean specialization surfaces this at the single-agent level.

The (PI)/Čencov observation. Two metric- α_2 cases (information-metric Kalman, Fisher-metric exp-family) are AAT-internally forced under the (PI) axiom in #scope-agent-identity combined with Čencov’s 1982 uniqueness theorem

— the fourth primary instance of the meta-pattern in #disc-additive-coordinate-forcing. The Euclidean metric- α_1 case is **AAT-internally derived via DA2'-inc equivalence** rather than theorem-imported, structurally connecting it to the bridge lemma in #form-composition-closure. The remaining metric- α_2 cases (Hessian, Lyapunov-linear-Hurwitz, Lyapunov-PID) are theorem-imported: the metric choice is problem-specific (matched to loss curvature or plant structure) rather than forced by an AAT-internal axiom.

Max attainable: *conditional*. The condition (the metric specification must be consistent with the system structure; the Slotine 2003 compositional theorems must apply to the topology) is inherent to the contraction-metric framework, not removable. Non-smooth systems (rule-based), strategic/game-theoretic equilibria, and systems without regular metric-state coupling remain honestly outside the template's scope (see §Honest failure modes).

Relationship to monotone-operator theory. The contraction-metric condition (CT2) is a specialization of monotone-operator theory's strong-monotonicity condition in a Hilbert space with a specific inner product structure (Rockafellar 1970; Bauschke & Combettes 2017 §§4, 22–28). AAT's #deriv-sector-condition reading of A2' α/β as operator-family classification — proximal / firmly nonexpansive / cocoercive / strongly-monotone-gradient / linear-PD — identifies the same mathematical content in operator-theoretic language; see that segment's Grounding paragraph for the correspondence. The present template does not claim the contraction-metric machinery as distinctively AAT — it claims the specific instantiation to AAT's state-variable choices and the composition with identifiability-floor + composition-scope-condition content as AAT-distinctive.

Honest failure modes.

The contraction-metric formulation does **not** lift AAT's coverage in five cases:

1. **Non-smooth dynamics.** Rule-based agents, state-machine controllers, threshold-triggered switches, and discontinuous policies fail the C^1 requirement. Piecewise-contraction extensions exist (Di Bernardo, Liuzza & Russo 2014, *SIAM J. Control Optim.* 52) but apply to specific switched structures, not general rule-based reasoning. This barrier is structural, not a gap in the metric machinery.
2. **Strategic / adversarial equilibria — three independent convergent obstructions.** The adversarial / strategic half of Part III is structurally outside contraction-metric scope for three distinct reasons that converge on the same conclusion; the convergence matters because it shows the limit is not an accident of one particular framework choice:
 - (i) *Compositional-contraction applicability* (Slotine 2003 §III): compositional contraction theorems require each sub-system to be individually contracting *and* for the composition topology to preserve contraction. Two-agent strategic dynamics with opposing objectives yield **saddle-point equilibria of the joint best-response dynamics** — attracting in some directions and repelling in others. Contraction requires attracting fixed points; saddle points break the framework directly.
 - (ii) *Passivity universality* (Khalil 2002 ch. 6 Thm 6.4): the dissipation inequality $\dot{S} \leq s(u, y)$ holds for *all* inputs u by the passivity property's universal quantifier. A strategic opponent can select u to maximize the recipient's storage-function growth rather than respecting the port structure. Passivity delivers cooperative/neutral robustness, not adversarial.
 - (iii) *Last-iterate convergence failure* (Daskalakis, Ilyas, Syrgkanis & Zeng 2018 “Training GANs with optimism” *ICLR*): adversarial gradient-ascent-descent dynamics (fictitious play, no-regret dynamics) converge to Nash equilibria *in average-iterate* sense (Cesàro-mean) but generically **fail last-iterate convergence** for non-zero-sum games. No contraction metric exists for systems whose last iterates do not converge to an attractor.

The three obstructions come from different theorem families (compositional-contraction / dissipativity-universality / differential-Lyapunov-attractor-theory) but converge on the same verdict: **contraction metrics cover the cooperative half of Part III** (#der-team-persistence, cooperative #deriv-critical-mass-composition, cooperative

#form-composition-closure). The adversarial / strategic half belongs to #deriv-strategic-composition (equilibrium convergence via Monderer-Shapley / Rosen) and #der-adversarial-destabilization (asymmetric attacker-target dynamics). This is a scope boundary that AAT correctly honors, not a gap to be closed — the apparatus that handles adversarial regimes is equilibrium-theoretic, structurally different from Lyapunov-contraction machinery.

3. **State-dependent metric stochastic coupling.** Pham & Slotine 2007 extends contraction to stochastic systems under small-noise bounds; the Itô-correction $c_{Itô}$ in (CT-S) is manageable for state-independent or slow-varying metrics. Strongly state-dependent M (adaptive-metric algorithms where M is learned alongside ξ) introduces cross-coupling that can destabilize contraction — the same structural issue #deriv-adaptive-gain-dynamics handles with its (MG-4) meta-gain coupling-boundedness condition.
4. **Multi-basin dynamics.** Lohmiller-Slotine contraction is local-on-a-region. Agents with multiple basins of attraction admit only basin-local metrics, giving a basin-chart structure rather than a single global metric. The basin-crossing boundary is exactly the #result-structural-adaptation-necessity trigger; the metric formulation makes the basin structure explicit but does not escape it.
5. **Identifiability-floor intersection.** The metric formulation operates downstream of identification. It assumes the correction function F points in a valid direction. Under the floors in #disc-identifiability-floor (on-policy L0-insufficiency no-go; L1' unobservable- C Cramér-Rao floor; composition-layer Instance 3 Liberzon no-go), the correction direction is structurally unavailable. Contraction metrics cannot recover identification — the two architectural moves are orthogonal.

Discussion.

Why factor this out. The Euclidean sector inequality in #result-sector-persistence-template (T2) matches a quadratic Lyapunov in Euclidean norm. For several AAT-relevant agent classes, the natural Lyapunov lives in a non-Euclidean metric (Fisher for statistical-manifold learning; Mahalanobis for Kalman; Hessian-induced for strongly-convex optimization; Lyapunov-equation-determined for asymmetric-stable linear systems). Stating AAT's persistence machinery once in parameter-free metric-formulation form makes visible that the sector condition and the contraction-metric condition are two states of the same apparatus — Euclidean and non-Euclidean — with the specific metric instantiated per sub-scope entry.

This is the *interior facet* of the stability certificate (5) in the non-Euclidean metric: the contraction-metric M is the certificate, and which metric is available per sub-scope is the certificate-strength ladder (R0-strict widest contraction / R1 cocoercive / R2 Čencov-forced; R0-loss is the wider Lyapunov-stable rung below where the certificate exists at $\kappa = 0$) of H . The five non-Euclidean metric cases are five certificates of one object, not five separate conditions.

The compositional payoff. The topology-indexed closures (CC-parallel / CC-cascade / CC-feedback) extend to heterogeneous-architecture composites via (CM2-M), closing the §6.1 heterogeneous-architecture obstruction in #deriv-critical-mass-composition. This is Part III's single largest coverage expansion: from “closed form for symmetric-matched-Tier-1” to “closed form under any of four topology classes with per-sub-agent (CT1)–(CT3) verification.”

Relationship to #result-sector-persistence-template. The two segments coexist. #result-sector-persistence-template's Euclidean formulation remains the default for agents whose natural coordinate is already Euclidean; #result-contraction-template is invoked when a non-Euclidean metric is natural (matrix Kalman, exp-family, ill-conditioned gradient, asymmetric-stable linear, PID-with-bounded-plant-nonlinearity). Every #result-sector-persistence-template invocation can be re-read as a #result-contraction-template invocation at $M = I$; the inverse mapping goes from metric to Euclidean with a $\kappa(M)$ degradation. Which direction is natural depends on the instance.

Relationship to #form-composition-closure's bridge lemma. The bridge lemma's DA2'-inc condition is equivalent to (CT2) with $M = I$ on convex C^1 domains. #result-contraction-template therefore sharpens the bridge-lemma understanding: Tier 1 (DA2'-inc globally) $\equiv$ metric-contracting-globally at $M = I$; Tier 2 (DA2'-inc locally with

$\kappa(D\delta)^2$ degradation) $\equiv$ metric-contracting-locally with a specific local metric; Tier 3 (per-domain verification) $\equiv$ no globally valid metric exists. The tier structure of the bridge lemma carries over to #result-contraction-template as Tier 1M / 2M / 3M respectively.

Relationship to #disc-identifiability-floor Instance 3 (composition-layer no-go). Instance 3's §3.3 construction uses matched-symmetric-Tier-1 scalar systems; the coupling-sign bit is unidentifiable from component marginals. Escape route (E-d) in that Instance names “shared metric structure (common contraction metric; Lohmiller-Slotine 1998)” as an adjacent machinery the no-go motivates; #result-contraction-template is the segment that machinery lives in. The Instance 3 escape is therefore operationalized by a #result-contraction-template invocation with a composite metric constructed from sub-agent metrics per the topology (parallel / cascade / feedback).

Relationship to meta-patterns.

- #disc-separability-pattern: metric- α_1 / metric- α_2 / metric- β is the seventh ladder (A2'-scope), closing the “binary special case” gap that segment's Discussion flagged. Separable-core (metric- α_1 , Euclidean, AAT-internally derived via DA2'-inc $\equiv$ (CT2)-M=I) / structured-repair (metric- α_2 , non-Euclidean, derived under explicit conditions including (PI)/Čencov for Fisher-metric cases) / general-open (metric- β , metric formulation fails).
- #disc-identifiability-floor: orthogonal axis. Contraction metrics lift the *geometry-of-correction* half of sub-scope α ; identifiability-floor operates on the *what-correction-to-make* half. The two architectural moves compose (Instance 3 at composition layer + metric-Tier structure at single-agent layer) but do not interact.
- #disc-additive-coordinate-forcing: the (PI)/Čencov axiom that grounds two of five metric- α_2 cases is the fourth primary instance of the meta-pattern, adding a Čencov-invariance machinery alongside the three Cauchy-FE instances. See that segment's Four Instances tables.

Findings.

Topology-Indexed Compositional Closures via Contraction-Metric Generalization. Brief: AAT's central machinery for proving an adaptive system holds its mismatch bounded was originally written in the simplest geometric setting (the equivalent of measuring distance with a ruler in flat space). Many real systems have a more natural geometry — a Kalman filter “thinks” in terms of an information ellipsoid, a learning agent thinks in Fisher-information geometry, a controller designed for an asymmetric plant thinks in a different metric still. This result restates the persistence machinery in a way that uses each system's natural geometry, lifting five additional agent classes from sub-scope β (assumed) to sub-scope α_1/α_2 (derived under explicit conditions), and inheriting Slotine 2003's topology-indexed compositional closures (parallel / cascade / negative-feedback) at the composition layer.

Impact: Part III's largest single coverage expansion: from “closed-form composite persistence under matched-symmetric Tier-1” to “closed-form composite persistence under any of four topology classes with per-sub-agent contraction verification, including heterogeneous-architecture mixes via (CM2-M).” For Part II, the structural-transparency observation that DA2'-inc $\equiv$ (CT2)-at-M = I shows AAT's existing Euclidean machinery was already carrying the Jacobian-level contraction commitment under another name; what changes is making that commitment legible and lifting non-Euclidean cases (Fisher-metric Kalman, exp-family natural-parameter, ill-conditioned strongly-convex gradient, asymmetric-stable linear, PID with bounded plant) from theorem-imported to AAT-internally derived under named axioms. Three independent obstructions (Slotine 2003 saddle-point; passivity-universality; Daskalakis 2018 last-iterate non-convergence) bound the template's scope honestly — adversarial composition is structurally outside contraction analysis and properly handled by #deriv-strategic-composition's equilibrium-convergence machinery.

Novelty Claim: *Claim synthesis* on contraction-metric machinery + AAT's sub-scope partition + (PI)/Čencov axiom. The contraction-metric primitive itself is standard nonlinear control theory; the topology-indexed closures are Slotine 2003. The synthesis is the AAT-internal mapping of metric choices to per-instance state-variable choices, the recognition that the Euclidean specialization is structurally the bridge-lemma's DA2'-inc, and the (PI)/Čencov axiom forcing two of the metric- α_2 cases to be AAT-internally derived rather than theorem-imported.

Related Work:

Table AI.1

<i>ASF concern</i>	<i>Prior-art language</i>	<i>Relationship / Positioning</i>
Differential-contraction primitive on Riemannian metrics	Lohmiller & Slotine 1998, “On Contraction Analysis for Non-linear Systems” <i>Automatica</i> 34:683–696 (published 1998, found pre-2026)	<i>formal antecedent</i> — supplies the (CT1)–(CT3) preconditions and the (CT-D) ultimate-bound argument adopted directly
Compositional contraction under topology	Slotine 2003, “Modular stability tools for distributed computation and control” <i>Int. J. Adapt. Control Signal Process.</i> 17:397–416 (published 2003, found pre-2026) — Theorems 1–3 (parallel; cascade; negative-feedback small-gain)	<i>formal antecedent</i> — supplies (CC-parallel) / (CC-cascade) / (CC-feedback) and the heterogeneous (CM2-M) extension. AAT adopts these directly, applied to the composition-validity question of #form-composition-closure
Finsler (direction-dependent) generalization	Forni & Sepulchre 2014, “A Differential Lyapunov Framework for Contraction Analysis” <i>IEEE TAC</i> 59 (published 2014, found 2026-04)	<i>adjacent literature</i> — Finsler-metric extension; flagged in Working Notes for highly anisotropic agents; not currently load-bearing
Strong monotonicity / firmly-nonexpansive operators	Rockafellar 1970, “On the maximality of sums of nonlinear monotone operators” <i>Trans. AMS</i> ; Bauschke & Combettes 2017, <i>Convex Analysis and Monotone Operator Theory in Hilbert Spaces</i> §§4, 22–28	<i>formal antecedent</i> — (CT2) at $M = I$ is strong monotonicity in Hilbert space; the Rockafellar-Wets / Nesterov equivalence (DA2’-inc $\Leftrightarrow$ Jacobian-symmetric-part PD $\Leftrightarrow$ (CT2)-at- $M = I$) is the standard monotone-operator identity AAT uses to make the Euclidean metric- α_1 case AAT-internally derived rather than theorem-imported
Fisher-metric / information-geometry uniqueness	Čencov 1982, <i>Statistical Decision Rules and Optimal Inference</i> AMS Translations 53	<i>formal antecedent</i> — under the (PI) axiom in #scope-agent-identity, Čencov’s theorem forces the Fisher-metric Kalman and exp-family natural-parameter cases AAT-internally
Bounded-loss composition under predictive compression	Subramanian, Sinha, Seraj & Mahajan 2020, “Approximate information state for approximate planning and reinforcement learning in partially observed systems” <i>arXiv:2010.08843</i> (published 2020, found 2026-04)	<i>adjacent literature</i> — closest mathematical neighbor to the bridge-lemma side of #form-composition-closure. The contraction-metric machinery in this segment supplies the (CT2) condition that, at $M = I$, is the bridge-lemma’s DA2’-inc; Sub20’s predictive-loss-to-control-error result lives in the same shape but at a coarser granularity (single-agent compressed control, not composite agent boundary)
Influence-based abstraction loss bounds	Congeduti, Mey & Oliehoek 2020, “Loss Bounds for Approximate Influence-Based Abstraction” <i>arXiv:2011.01788</i> (published 2020, found 2026-04)	<i>adjacent literature</i> — bounds value loss from approximate cross-agent influence; the (CC-feedback) topology with bounded loop gain is the contraction-metric counterpart at the dynamics layer
Saddle-point / adversarial limits of contraction	Daskalakis, Ilyas, Syrgkanis & Zeng 2018, “Training GANs with optimism” <i>ICLR</i> (published 2018, found 2026-04); Khalil 2002 <i>Nonlinear Systems</i> 3rd ed. ch. 6 Thm 6.4 (passivity-universality)	<i>contradicted by this finding</i> (in the sense of correctly bounding scope) — three independent obstructions to handling adversarial regimes converge: Slotine 2003 §III saddle-point limitation; passivity-universality’s universal quantifier; Daskalakis 2018 last-iterate non-convergence. The convergence demonstrates the scope boundary is not an artifact of one framework; adversarial composition is structurally outside contraction analysis

<i>ASF concern</i>	<i>Prior-art language</i>	<i>Relationship / Positioning</i>
Switched-systems common-Lyapunov nonexistence	Liberzon 2003, <i>Switching in Systems and Control</i> Theorem 2.1	<i>formal antecedent</i> — supplies #disc-identifiability-floor Instance 3's no-go; the contraction-metric construction operationalizes escape route (d) (common contraction metric) for that Instance

Search Log:

- 2026-04 (*nominally comprehensive*): The composition-layer extension of this segment (the (CC-*) closures and (CM2-M)) was reviewed under the same Undermind pass that covered #form-composition-closure. Verdict: the contraction-metric machinery itself is standard adopted prior art (formal antecedent); the AAT-internal mapping of metric choices to per-instance sub-scope α_1/α_2 entries and the use of (PI)/Čencov to force two metric- α_2 cases AAT-internally are the synthesis-level contributions. The Sub20 / Con20 positioning at the composition-validity layer is recorded both here and at #form-composition-closure since the contraction template is what supplies the formal mechanism the bridge lemma uses.
- 2026-04 (*targeted*): Lohmiller-Slotine 1998 / Slotine 2003 / Forni-Sepulchre 2014 known and cited from the contraction-analysis literature. The three adversarial obstructions (Slotine 2003 §III; Khalil 2002 ch. 6; Daskalakis et al. 2018) are convergent scope-boundary evidence.
- 2026-04 (*intuition-only*, prior): expected the contraction-metric primitive to be standard control-theory; expected the composition theorems to be the more proprietary content. Comprehensive search confirmed both: the primitive is well-established, and the composition theorems are Slotine 2003's, with AAT's contribution lying in the application to #form-composition-closure's admissibility-and-bridge-lemma framework.

Working Notes.

- **Landing context.** The segment is the primary result of spikes/.integrated/spike-contraction-metric-generalization.md (2026-04-23 Gap A/B cycle). The Jacobian-level observation (DA2'-inc $\equiv$ (CT2)-M=I) and the (PI)/Čencov observation derive from the jacobian-b1 strengthening (2026-04-23 follow-up; absorbed in the spike-routing cycle, CHANGELOG 2026-05-17).
- **Heterogeneous general-graph composition.** Slotine 2003 §IV's virtual-system technique covers general directed graphs under a small-gain-along-every-cycle condition. Operationalizing this for #form-composition-closure's general multi-agent setting — including identifying which cycle-gains to measure in practice — is a follow-up derivation. (CM2-M) in this segment handles 2-agent negative-feedback; $N \geq 3$ with general topology remains open but structurally clear.
- **Jacobian-level B1 strengthening is open.** The (CT2) condition on F 's Jacobian is a Jacobian-level B1 analog. Whether the integral B1 of #der-gain-sector-bridge can be *strengthened* to a Jacobian-level B1* that implies (CT2) AAT-internally (rather than via theorem-import from Lohmiller-Slotine) was attacked in the jacobian-b1 strengthening (spike-routing cycle, CHANGELOG 2026-05-17). Three of the five metric- α_2 cases (Hessian, Lyapunov-linear-Hurwitz, Lyapunov-PID) remain theorem-imported after that spike; two (Fisher-metric-Kalman, Fisher-metric-exp-family) lift AAT-internally under (PI)/Čencov. The remaining gap is the three theorem-imported cases.
- **Candidate strong-option promotion.** Under a *heredity* axiom (composite admissibility derivable from sub-agent properties — a legitimate AAT-internal strengthening of #disc-composition-consistency), agent-level B1* is forced, metric-Tier structure collapses, and (CM2-M) promotes from Slotine-imported to AAT-derived. Flagged in the Jacobian-level B1 spike's §II strong option; promotion is an architectural decision requiring its own scoping.
- **Finsler metrics.** Forni & Sepulchre 2014 extends the differential Lyapunov framework to Finsler (direction-dependent) metrics. Likely useful only for highly anisotropic corrections (coordinate-asymmetric agents); not priority.

- **The (T2) / (CT2) structural observation matters for Part I too.** #result-persistence-condition's linear special case ($F = \mathcal{T}\delta$, $\alpha = \mathcal{T}$, (T2) global) is trivially contraction-metric with $M = I$ and $\lambda = \alpha$. But non-linear #result-persistence-condition instances (nonlinear mismatch dynamics, gradient on strongly convex losses) would benefit from contraction-metric framing to remove the Euclidean-norm-specific condition-number penalty in the persistence constant. A cross-reference from #result-persistence-condition back to this segment may be warranted after the next promotion pass.
-

Appendix AJ

Derivation: Variational Approximate-A2' (ε -Fidelity)

Variational / approximate-posterior agents (VI, amortized VI, active-inference-style variational free energy) currently sit in A2' sub-scope β per #form-sector-condition: their correction functions target the *best-in-class* variational posterior q^* rather than the true posterior p , and the approximation gap can rotate the correction direction enough to break B1 directional fidelity (4.3). Under a KL bound $\text{KL}(q_\phi \| p) \leq \varepsilon$ on the variational approximation, directional fidelity recovers in a quantifiable form: **ε -fidelity B1**, with sector-constant degradation scaling as $O(\sqrt{\varepsilon})$ (Pinsker-tight). The sector-persistence template applies under a Regime-A / Regime-B decomposition — clean sector bound on an annulus away from the projection-error floor, approximation-dominated on a ball of radius $\delta_0 = O(\sqrt{\varepsilon})$ around the target. This promotes controlled-KL VI from sub-scope β to a new intermediate tier α' within the A2' partition (cf. the $\alpha / \alpha_1 / \alpha_2 / \beta$ refinements from #deriv-adaptive-gain-dynamics and #deriv-fisher-whitened-update-rule).

The result has a striking operational consequence: **mean-field VI is structurally sub-optimal as an adaptive system by a factor $\sqrt{1 - r^2}$ in the sector constant, regardless of compute or training data**, because the mean-field family does not match the Fisher metric. **Natural-gradient VI recovers the full sector parameter** (Khan-Lin 2017 conjugate-computation equivalence places it in full sub-scope α rather than α') because the natural-gradient direction matches the Fisher metric — no degradation. Where persistence-optimality matters, the variational-family choice is not interchangeable; the long-running Bayesian-deep-learning practical preference for natural-gradient over mean-field has a principled persistence-theoretic reading inside AAT.

Formal Expression.

ε -fidelity B1.

Let the true posterior be $p(z | x)$ and the variational approximation $q_\phi(z | x)$ with $\text{KL}(q_\phi \| p) \leq \varepsilon$. Under standard Lipschitz assumptions on the observation model and nested-support on the variational family, the total-variation distance bounds as $\|q_\phi - p\|_{TV} \leq \sqrt{\varepsilon/2}$ (Pinsker's inequality). Propagating this bound through the correction function via Cauchy-Schwarz:

Derived (eps-fidelity-B1, from Pinsker + Cauchy-Schwarz).

$$(\hat{K}\hat{P} - P^*)^T(\hat{P} - P^*) \geq (c_{\min} - C_H\sqrt{2\varepsilon/\|\delta\|}) \cdot \|\delta\|^2 \quad (\text{AJ.1})$$

where C_H is a constant depending on the observation-model Lipschitz constant. The effective sector constant

$$c_\varepsilon(\|\delta\|) = c_{\min} - C_H \sqrt{2\varepsilon/\|\delta\|} \quad (\text{AJ.2})$$

is state-dependent: large-mismatch regions see near-full sector constant; near-target regions see degraded sector constant due to approximation error dominating.

Regime-A / Regime-B decomposition.

Define $\delta_0 = 2C_H \sqrt{2\varepsilon/c_{\min}}$ (the approximation-dominated radius around target). Then:

Formulation (regime-decomposition-variational).

- **Regime A — clean sector bound** ($\|\delta\| > 2\delta_0$). On the annulus $\mathcal{B}_R \setminus \mathcal{B}_{2\delta_0}$, $c_\varepsilon \geq c_{\min}/2$: sector condition holds with constant $c_{\min}/2$. The template's ultimate bound $\rho_\xi/(c_{\min}/2)$ applies.
- **Regime B — approximation-dominated floor** ($\|\delta\| \leq 2\delta_0$). Near target, c_ε can be arbitrarily small; the correction may not contract. The ultimate bound gains an additive $O(\sqrt{\varepsilon})$ term: $R_\varepsilon^* = \rho_\xi/(c_{\min}/2) + \delta_0 = \rho_\xi/(c_{\min}/2) + O(\sqrt{\varepsilon})$.

Sector-persistence template instantiates with: - State variable $\xi = \hat{P} - P^*$ (variational mismatch to target posterior). - Effective ultimate bound R_ε^* on Euclidean norm. - Persistence requires $R_\varepsilon^* < R$, i.e., $\rho_\xi/(c_{\min}/2) + O(\sqrt{\varepsilon}) < R$.

Khasminskii stopping-time localization (same technique as #deriv-sector-condition Prop A.1S and the A2' strengthening spike) applies to the annulus Regime A; Regime B is handled by accepting δ_0 as a projection-error floor.

Per-case verdicts.

The α' -membership depends on the specific variational scheme:

Derived (per-variational-case).

Natural-gradient VI with exponential-family q_ϕ . Khan & Lin 2017 (*Conjugate-computation variational inference*) showed natural-gradient VI is equivalent to closed-form conjugate-Bayesian updates for exponential-family variational distributions. This recovers *full* sub-scope α membership (not merely α'), with A2' derived rather than ε -degraded, by converting the variational update into a Bayesian update in a reparameterized family. The $\varepsilon = 0$ limit is exact.

Mean-field VI. When the true posterior is approximately factorized ($p(z) \approx \prod_i p_i(z_i)$), mean-field VI achieves small ε and is the workhorse α' case. Ultimate bound degrades by additive $O(\sqrt{\varepsilon})$; sector persistence holds.

Amortized VI (VAE-style). Amortization adds a second approximation error (the variational network's function-class limit). KL bounds compose *additively*: δ_0 grows as $\sqrt{\varepsilon_{\text{family}} + \varepsilon_{\text{amort}} + \varepsilon_{\text{generalization}}}$. Under controlled- ε amortized VI, α' membership holds with a larger floor.

Diffusion-posterior / energy-based with uncontrolled MCMC. No controlled ε bound; ε grows with mixing time. Stays firmly in sub-scope β .

Active inference (variational free energy). Conditional α' under exponential-family q + natural-gradient; ε -degraded α' otherwise. This does **not** force V-strong G-BP2 (presentation of AAT as control-theoretic specialization of active inference) — V-medium (KL-divergence-based cognitive cost) remains the appropriate scope commitment; the comparison trail is in the spike-routing cycle (CHANGELOG 2026-05-17).

Sub-scope α' in the A2' partition.

The A2' sub-scope partition is structured as:

Derived (sub-scope-alpha-prime-partition).

- **Sub-scope α** (derived under B1 directional fidelity; #der-gain-sector-bridge): Kalman, conjugate Bayesian, exponential-family natural parameters, strongly-convex gradient, L2-regularized, linear-PD-symmetric.
- **Sub-scope α_1/α_2** (#result-contraction-template): metric-formulation generalization of α .
- **Sub-scope α_3** (#deriv-fisher-whitened-update-rule): correlated evidence + Fisher-whitened under (PI)/Čen-cov.
- **Sub-scope α'** (this segment): controlled-KL VI under Pinsker's inequality with $O(\sqrt{\varepsilon})$ sector-constant degradation + Regime-A/B decomposition.
- **Sub-scope β** : uncontrolled- ε agents; non-smooth rule-based; severely misspecified; per-step SGD; human judgment.

This gives the full current picture: $\{\alpha, \alpha_1, \alpha_2, \alpha_3, \alpha', \beta\}$.

Epistemic Status.

Conditional. Max attainable: *exact* for the Pinsker-bound derivation (standard inequality). The $O(\sqrt{\varepsilon})$ rate is tight in general — improving to $O(\varepsilon)$ requires stronger inequalities (log-Sobolev, Talagrand T2) holding for the specific variational family and target; not chased in this segment but is a potential refinement for concentration-preserving families.

Load-bearing: - Pinsker + Cauchy-Schwarz gives the $O(\sqrt{\varepsilon})$ rate rigorously. - Regime-A annulus sector bound with $c_{\min}/2$ is derived (straightforward bookkeeping on the state-dependent c_ε). - Natural-gradient VI promotion to full α is standard (Khan & Lin 2017). - Mean-field VI as workhorse α' is standard variational analysis.

Not established: - $O(\varepsilon)$ rate (would require stronger inequalities). - Quantitative ε bounds for specific variational families in practice (amortization gaps are typically not estimable tightly). - α' membership under mode-missing VI (multimodal posteriors where q concentrates on a subset of modes). These may require separate analysis.

Honest Failure Modes.

- **Uncontrolled- ε agents:** diffusion posteriors with finite MCMC; VAEs trained without KL control; any VI without a bounded- ε certificate. Stays in β .

Scope honesty — variational-alpha-prime-limits.

- **Mode-missing on strongly multimodal posteriors:** q misses modes; ε is irreducible regardless of optimization. α' fails structurally.
- **Non-Lipschitz observation models:** the C_H constant diverges; the Pinsker propagation bound doesn't hold. α' fails.
- **Support mismatch:** q 's support is not a subset of p 's support (or vice versa); KL is $+\infty$. α' fails.
- **Non-exponential-family + non-amortized + uncontrolled + multimodal:** the worst case, strictly β .

Discussion.

Relationship to #form-strategy-complexity-cost. The G-BP2 V-medium variational form landed in #form-strategy-complexity-cost uses KL divergence in the cognitive-cost term. This segment provides the complementary story on the persistence side: KL-bounded VI has ε -fidelity B1 and α' sector condition. The two together characterize the cognitive/persistence tradeoff for variational agents: ε controls both the cognitive cost (how far from the target) and the persistence degradation (how much sector-constant penalty). Large ε means cheap-but-persistently-weak; small ε means expensive-but-persistently-sharp.

Relationship to #disc-compression-operations. Variational compression is the first of the four AAT compression operations; its α' sector structure gives a concrete persistence guarantee for variational representation of M_t . The other three operations (strategy Σ_t , shared-intent, coarse-graining Λ) may admit similar ε -fidelity analyses where a KL-bound on the compression is available.

Meta-pattern positioning. - #disc-separability-pattern: α' sits on the structured-repair tier of the A2'-scope ladder (7th ladder from #result-contraction-template), alongside metric- α_2 and α_3 . Each represents "sector condition recovered under explicit additional structure." - #disc-additive-coordinate-forcing: variational persistence sits *outside* this meta-pattern — no logarithmic coordinate is forced; Pinsker is a bound, not a Cauchy-FE argument. α' is an adjacent family member (shape: controlled-approximation-with-quantified-degradation), not a primary instance.

Working Notes.

- Landing context: spikes/.integrated/spike-variational-a2prime.md (2026-04-23 Gap A/B cycle).
 - **Open: $O(\varepsilon)$ rate via Talagrand T2.** For concentration-preserving variational families with log-Sobolev or T2-Talagrand inequalities, the Pinsker bound can be sharpened to $O(\varepsilon)$ directly (without square-root). Not chased here; would extend α' to a tighter sub-tier for specific family classes.
 - **Amortization gap bounds.** Generalization-gap analysis for VAEs (Rainforth et al. 2018; Nowozin 2018) gives function-class error bounds; composing with variational-family error gives the full δ_0 in amortized VI. Follow-up work.
 - **Connection to natural-gradient VI.** Khan & Lin 2017's conjugate-computation framing makes natural-gradient VI full sub-scope α (not just α'). The α' tier is for *non-natural-gradient* VI under KL control.
-

Appendix AK

Derivation: L1' Update-Bias Formula

Under L1' correlated-evidence regimes with unobservable common cause, the default log-odds edge-update (U) converges to a **biased fixed point** — edges settle at log-odds values that match the L1' root-probability rather than the marginal-edge truth. This segment derives a closed-form bias formula for the two-edge OR-root case, verifies it via Monte Carlo, and composes it with #disc-identifiability-floor Instance 2 and #schema-strategy-persistence's forgetting prerequisite to produce a **dual forgetting-rate requirement** on strategic persistence.

The bias is the quantitative companion to F13's structural-identifiability-floor result: F13 shows identification is *structurally impossible* from single-channel observation; this segment shows the resulting *magnitude drift* is bounded, Lipschitz in correlation strength ρ , and Cramér-Rao-floored under forgetting. The dual forgetting-rate requirement reads operationally: forget *fast enough* to track non-stationarity (the plasticity lower bound from 9.10) *and* fast enough to keep the L1' bias bounded (this segment's upper-bound-driving lower bound on $1 - \lambda$). When the two compose with #form-consolidation-dynamics's stability cadence into an empty admissibility window, the agent must augment its observability or accept biased-fixed-point operation — the operational signature of a forgetting-window infeasibility in which constraint is binding.

Formal Expression.

Closed-form bias at matched-marginal initial conditions (two-edge OR-root).

Let root node G have two direct causal parents e_1, e_2 via OR-aggregation under L1' with latent common cause C producing correlation ρ between the parent edges. Assume matched-marginal initial conditions: $\hat{\mu}_j^{\text{init}} = \mu_j^*$ for each edge's marginal probability. Under the default log-odds update with per-edge gain $\eta_{\text{edge}} = 1/(n_k + 1)$ and identifiability coefficient ι_k , the per-cycle bias on edge k 's log-odds update satisfies

Derived (L1-bias-formula-OR-root).

$$B_k(\rho) = -\frac{\iota_k \cdot (1 - \mu_{\bar{k}}) \cdot \rho}{(n_k + 1) \cdot [(1 - \mu_1)^2 + (1 - \mu_2)^2]} \quad (\text{AK.1})$$

where $\bar{k} \in \{1, 2\} \setminus \{k\}$ is the sibling edge. The bias is linear in ρ , decays as $1/(n_k + 1)$ with experience, topologically asymmetric via the Jacobian factor $(1 - \mu_{\bar{k}})/\|\mathbf{J}\|^2$.

Admissible region (precise). The denominator is the squared OR-aggregation Jacobian norm $\|\mathbf{J}\|^2 = (1 - \mu_1)^2 + (1 - \mu_2)^2$. The first-order perturbation is well-posed precisely when $\|\mathbf{J}\|$ is bounded below by some fixed $\epsilon_J > 0$ — equivalently, when at least one edge probability is bounded away from 1. This is the **Jacobian-norm admissibility**

condition $\|\mathbf{J}\| \geq \epsilon_J$; it is *weaker* than requiring both μ_1, μ_2 to be bounded away from 1, and it is *sharper* than the qualitative “bounded away from determinism” reading: only the joint deterministic-success corner $\mu_1, \mu_2 \rightarrow 1$ together is excluded. As either edge alone approaches deterministic success the bias formula remains valid (the surviving sibling edge carries the Jacobian norm); only the diagonal corner where both edges saturate is excluded, and there the per-cycle update itself vanishes (observations carry no residual information), so “no learning, no bias” is the operationally correct behaviour and the formula’s apparent singularity is degenerate rather than physical. Inside the admissible region $\|\mathbf{J}\| \geq \epsilon_J$, the closed-form predictions are quantitatively verified by Monte Carlo; at the diagonal corner the segment defers to #deriv-edge-credence-dynamics’s degenerate-update analysis. The Jacobian-norm condition is the natural admissibility statement because the bias formula is the projection of the update onto $\mathbf{J}/\|\mathbf{J}\|^2$, so well-posedness requires only that this projection be defined.

For AND-root aggregation, the bias has the same structural form with opposite sign. For mixed OR-AND sub-plans, the bias decomposes sub-plan-wise.

L1'-induced biased fixed point.

Under continued L1' evidence, the log-odds fixed point satisfies

Derived (biased-fixed-point-exists).

$$\lambda_k^{\text{L1'-FP}} = \lambda_k^* + \int_0^\infty e^{-\eta_{\text{edge}} t_k t} B_k(\rho) dt = \lambda_k^* - \rho \cdot (1 - \mu_k) / [(1 - \mu_1)^2 + (1 - \mu_2)^2] \quad (\text{AK.2})$$

(for constant- t approximation). Each edge’s log-odds converges to a **wrong value** that matches the L1' root probability rather than the marginal truth. **Plan-level correctness is purchased at edge-level miscalibration:** the agent predicts plan outcomes correctly on-policy (the L1' mixture on the sibling matches the L1' marginal), but its per-edge credences are systematically biased.

This is the quantitative companion to F1’s on-policy undetectability (#disc-identifiability-floor Instance 1) at the marginal level: the bias accumulates at rate $1/(n_k + 1)$, converges to a bounded non-zero value, and is *not detectable from on-policy data* because the mixture’s marginal is matched to the on-policy observations.

Observable-C zero-bias result.

Under Prop B.7’s five-way-gating with C observable, conditioning on C decomposes the L1' mixture into two single-child-per-component problems:

Derived (observable-C-zero-bias).

$$B_k(\rho \mid C = c) = 0 \quad \text{in expectation at conditional truth, for each } c \in \{0, 1\}. \quad (\text{AK.3})$$

Prop B.7’s decomposition *exactly eliminates the bias* — the five gating conditions (facilitator monotonicity; C -observability; $t > 0$ for common-cause edge; marginal sector condition; identifiability of per-component rates) are precisely the conditions under which conditional-on- C updates restore exactness. This makes Prop B.7’s derivation the quantitative escape from the L1' bias.

Unobservable-C Cramér-Rao bias floor under forgetting.

When C is unobservable and the agent uses experience discounting at rate $\lambda \in [0, 1)$ to maintain plasticity (#schema-strategy-persistence's forgetting prerequisite), the steady-state bias is

Derived (bias-floor-under-forgetting).

$$B_k^{SS}(\rho, \lambda) = -\frac{\iota_k \cdot \rho \cdot (1 - \mu_{\bar{k}})}{(1 - \lambda) \cdot [(1 - \mu_1)^2 + (1 - \mu_2)^2]}. \quad (\text{AK.4})$$

The Cramér-Rao floor from #disc-identifiability-floor Instance 2's rank-1 Fisher matrix translates directly: no unbiased estimator can reduce the bias below $\rho/(1 - \lambda) \cdot [\text{constant}]$. The bias is structurally present in any online estimator operating on single-channel observation of one child edge, not merely the specific default-signal estimator.

Dual forgetting-rate requirement.

#schema-strategy-persistence requires forgetting rate $(1 - \lambda) > \rho_{\Sigma}/R_{\Sigma}$ for asymptotic sector-persistence. The L1' bias floor above imposes a **second** constraint: to keep the bias bounded relative to the sector reserve,

Derived (dual-forgetting-requirement).

$$(1 - \lambda) > c_B \cdot \rho/R_{\Sigma}^{\text{bias}} \quad (\text{AK.5})$$

where c_B is a topology-dependent constant from the Jacobian factor and R_{Σ}^{bias} is the tolerable bias-reserve within the sector region. The forgetting rate must satisfy **both**:

$$(1 - \lambda) > \max\left(\frac{\rho_{\Sigma}}{R_{\Sigma}}, \frac{c_B \cdot \rho}{R_{\Sigma}^{\text{bias}}}\right). \quad (\text{AK.6})$$

When the admissibility window is empty (no forgetting rate satisfies both constraints simultaneously), no standard persistence regime works — the agent must augment (observable-C via instrumented regime indicators; multi-child joint observation; plan-level fallback) or accept biased-fixed-point operation.

Monte Carlo verification.

Two independent Monte Carlo implementations confirm the closed-form predictions — the original 2026-04-23 record (four scenarios: symmetric strong common-cause, asymmetric mild, degenerate $\rho = 0$, θ_C sweep; 400 trials $\times$ 5000 cycles, OR-root) and a fresh 2026-07-16 implementation that reproduces the original's scenario-A cumulative drift to 4 significant figures (-0.3996 vs -0.3997) and extends verification to the AND-root case: - Initial-cycle rate matches the closed form within $< 1\%$ for **both OR and AND roots** across the feasible covariance range $\rho \in \{0.04, 0.09, 0.16, 0.20, 0.24\}$ (note $\rho = \text{Cov}(Y_1, Y_2) \leq 1/4$ by construction) and in the asymmetric-edge case. - Sign matches prediction everywhere; **OR and AND roots produce opposite-sign biases of matching magnitude** (AND-root now simulated, not derivation-only). - Vanishing at $\rho = 0$ verified (degenerate L1'); no false positive under L0-matched truth (drift $\approx 3 \times 10^{-4}$ over 5000 cycles). - Cumulative drift is log-order (far from linear) and decelerates slightly faster than the frozen-state $\sum_t 1/(n + t + 1)$ prediction, consistent with the bias-induced $\mathbf{p}(t)$ trajectory effect; the exact transient shape has no closed form (see Epistemic Status).

Empirical claim (monte-carlo-confirmation).

Epistemic Status.

Conditional. Max attainable: *exact* for the two-edge OR-root / AND-root matched-marginal case. The closed-form bias formula is derived via first-order perturbation of the log-odds update around matched-marginal truth; the Jacobian factor is specific to the DAG topology (OR vs AND vs mixed). *Robust qualitative* for deeper DAG structures where the bias Decompose sub-plan-wise but closed-form coefficients depend on per-sub-plan Jacobians.

What is load-bearing: - The linear-in- ρ scaling is exact (first-order in small correlation). - The $1/(n_k + 1)$ experience scaling is exact (follows from #deriv-edge-update-natural-parameter's Beta-Bernoulli gain). - The observable- C zero-bias result under Prop B.7 gating is exact. - The dual forgetting-rate requirement is a first-order composition of two separate constraints.

What is not established: - Closed-form bias for non-matched-marginal initial conditions (transient bias during convergence). Numerical simulation works; analytic form is messier. - Higher-than-second-order DAG structures (L2-general). Scope boundary. - Full coupled bias formula under #deriv-adaptive-gain-dynamics meta-gain machinery. Open.

Discussion.

Composition with #disc-identifiability-floor Instance 2. F13 establishes the structural no-go for single-channel L1' mixture identification (Cramér-Rao rank-1); this segment gives the **quantitative numerical-floor** companion — the bias is bounded, continuous in ρ , and vanishes at $\rho = 0$. Instance 2 answers “is it identifiable?” (no, structurally); this segment answers “if we update anyway, how wrong is the result?” (bounded drift to a biased fixed point).

Composition with #deriv-fisher-whitened-update-rule. The Fisher-whitened update recovers sharp *directional* fidelity but does not reduce the *magnitude* bias — Fisher whitening operates downstream of identification, and Instance 2's obstruction is an identification obstruction. The full picture for the default signal function under L1': **direction preserved (angle $\leq 45^\circ$ per #deriv-fisher-whitened-update-rule), magnitude biased by $B_k(\rho)$ with Cramér-Rao floor under forgetting (this segment), observable- C restores exactness (Prop B.7), undetectable on-policy (Instance 1 + Instance 2 composition).** This is the honest quantitative account of what the default signal function does under Gemini's Gap A (validation under correlated failures).

Relationship to #schema-strategy-persistence's forgetting. The segment's existing forgetting prerequisite (for asymptotic sector-persistence) becomes dual under L1': forgetting must *both* outpace sector disturbance *and* prevent bias accumulation. When the two constraints are compatible (admissibility window non-empty), the existing forgetting-rate requirement suffices with appropriate ρ -dependent tuning. When incompatible, augmentation is required — a structural scope narrowing that elevates observable- C instrumentation from convenience to prerequisite.

Connection to X. The detection-latency theorem ($\mathbb{E}[T_{\text{detect}}] = \Omega((n_{\min} + 1)/\varepsilon)$) assumes unbiased marginal tracking; under L1' bias, the pre-accumulated drift degrades effective ε for subsequent regime-change detection. High-operating-point agents with pre-accumulated L1' bias face a compounded myopia floor. Follow-on refinement possible in X.

Working Notes.

- Landing context / reproducibility: the closed-form derivation and the original Monte Carlo record (scenarios A–D, OR-root) are in spikes/.integrated/spike-l1-update-bias.md §7 (2026-04-23 Gap A/B cycle; its script lived in /tmp and is lost). The 2026-07-16 re-verification — fresh implementation, full script preserved, OR+AND roots, feasible- ρ grid, 4-sig-fig reproduction of the original scenario-A drift — is spikes/spike-l1-bias-sim-rerun-2026-07-16.md. Citable-supplement release at publication should use the rerun's script.
- 2026-07-16 — Monte Carlo block corrected (truth-status repair from the bulk-64 verify-pass). The block previously described a simulation that was never run (“four scenarios: OR/AND $\times$ cooperative/adversarial”; “ $\rho \in \{0.1, \dots, 0.9\}$ ” — infeasible, since $\text{Cov} \leq 1/4$ under the spike's parameterization): a confabulated summary written at landing, while the

actual 2026-04-23 record was sound (it reproduces exactly). Restated to what ran, strengthened by the rerun (AND-root now verified). Guard: when summarizing a simulation record into a segment, transcribe its parameters — don't paraphrase from memory.

- **Open: non-matched-marginal transient bias.** Closed-form for arbitrary initial conditions is messier but numerically tractable. A tighter analytic treatment would compose `#deriv-edge-update-natural-parameter's` β -Bernoulli dynamics with the L1' mixture and likely requires Markov-chain convergence-rate analysis.
 - **Open: N -edge common-cause extensions.** Three or more children of a single common cause produce a tensor of Jacobian factors. The structural scaling likely remains $O(\rho)$ per edge, but the exact coefficients need working out.
-

Appendix AL

Derived: Architecture No-Identifiability from On-Policy Summary Data

The framework’s identifiability-floor at the *agent-internal architecture* layer. Two AAT agents whose linearized closed-loop residual dynamics are minimal state-space realizations related by an invertible similarity transformation produce *identical* on-policy summary statistics: identical innovation spectrum, identical sector-condition summary (α, R) , identical similarity-invariant deviation summaries — and therefore identical observation law under on-policy, in-regime, summary-only access. An observer with that access cannot distinguish the two agents, even though they implement the correction in structurally distinct internal coordinates. The architectural degree of freedom AAT cares about — #der-agent-opacity’s “which internal mechanism” — lives entirely in the $GL(n)$ similarity fiber and is annihilated by the on-policy observation map.

The no-go is *exact* in the linear-Gaussian / Kalman sub-scope, anchored by Kalman 1963 and Ho-Kalman 1966 canonical-form non-uniqueness; *robust-qualitative* in the general sub-scope via the Causal Hierarchy Theorem applied at the agent-as-SCM layer (Bareinboim, Correa, Ibeling & Icard 2022). The mechanism is **not** Sylvester’s law of inertia inherited for free — the *generating* group action that produces the indistinguishable pair is state-space similarity $T(\cdot)T^{-1}$, not the metric congruence $S^T(\cdot)S$ that Sylvester governs. It *reduces*, however, to the same shape as the Instance-2 rank-collapse floor (R Prop B.7): the on-policy Fisher information of the realization manifold is identically zero along the similarity-fiber tangent distribution, so the observed-Fisher is rank-deficient on a structurally-forced indeterminacy manifold — Instance-2’s mechanism with the manifold being a Lie-group fiber rather than a single indeterminacy sub-manifold. Escape-irreducibility under the agent’s *metric/coordinate* freedom (distinct from the similarity freedom that *generates* the pair) is then Sylvester at one remove, identically to Instance 2.

Three structurally distinct escapes: (a) loop-interventional access (7.3) perturbs along similarity-fiber directions that on-policy data never excites — the intervention generates Level-2 data, the no-go is Level-1-only; (b) higher-moment / out-of-regime observation distinguishes the pair in the nonlinear sub-scope (AAT’s β) — *provably void* in the linear-Gaussian sub-scope, since a Gaussian innovation carries no information beyond second order; (c) architecture instrumentation (direct read of the update rule or internal state) breaks the black-box scope, structurally the same kind as Instance 2’s “instrument the latent” escape. A fourth proposed escape — horizon extension under the same policy — collapses into (a): more samples of an annihilated observation law cannot escape the annihilation.

The strengthened consequence is that #der-loop-interventional-access acquires a *third semantically-distinct deployment mode* — at the agent-internal layer — alongside Instance 1’s agent-self-intervention (breaking causal-sufficiency degeneracy) and Instance 3’s observer-on-sub-agent intervention (breaking coupling-sign degeneracy). The shared content across the three modes is “Level-2 data breaks a Level-1 degeneracy”; the distinct content is *which* degeneracy.

Formal Expression.

The Setting.

Two AAT agents A and A' are *architecturally distinct, behaviorally identical from on-policy summary data* in a regime Ω when:

Definition (architecture-noidentifiability-setting).

- both operate on the same regime-restricted trajectory under their respective on-policy execution;
- both share the same sector-condition summary (α, R) in $\Omega(A)$, equivalently the same ultimate-deviation statistics on the similarity-invariant summaries of $\|\delta\|$ in that regime;
- the observer has *on-policy, in-regime, summary-only* access — no off-policy intervention, no out-of-regime samples, no white-box read of the update rule or internal state.

Per #scope-agent-identity, A and A' are agents on singular trajectories with action-marginal invariance under change of internal-coordinate parameterization. The question is whether the architectural content — the basis in which the correction is implemented, internal to the agent — can be recovered from the observer’s accessible data.

The No-Go Theorem (exact in the linear-Gaussian / Kalman sub-scope).

Let the linearized closed-loop residual dynamics of A and A' be minimal state-space realizations

Derived (architecture-noidentifiability-Kalman, from Kalman-Ho canonical-form non-uniqueness applied to AAT-agent realizations), exact in the linear-Gaussian sub-scope.

$$d\delta_t = -F \delta_t dt + \sigma_w dW_t, \quad d\delta'_t = -F' \delta'_t dt + \sigma'_w dW_t \quad (\text{AL.1})$$

with F, F' Hurwitz (sector-condition-satisfying per #deriv-sector-condition) and the two realizations related by an invertible similarity $T \in GL(n)$:

$$F' = T F T^{-1}, \quad \sigma'_w \sigma'^{\top}_w = T \sigma_w \sigma_w^{\top} T^{\top}. \quad (\text{AL.2})$$

Claim. The stationary innovation process — hence the (α, R) -summary, the innovation spectrum, and every *similarity-invariant* summary of the on-policy observation law — is **identical** for A and A' . No on-policy, in-regime, summary-restricted statistic distinguishes them.

Derivation. The innovation’s stationary law is determined by the pair $(F, \sigma_w \sigma_w^{\top})$ only through similarity-invariant quantities. Specifically:

- The eigenvalues of F (the α -spectrum entering the sector summary) are similarity-invariant.
- The stationary covariance Π solves the Lyapunov equation

$$F \Pi + \Pi F^{\top} = \sigma_w \sigma_w^{\top} \quad (\text{AL.3})$$

(standard form for the OU drift $-F$ with F Hurwitz, sign convention matching #deriv-sector-condition), and transforms under similarity as $\Pi' = T \Pi T^{\top}$.

- The observed quantities are the *innovation spectrum* $\det(i\omega I - F)$ and its similarity-invariant moments — for example, the *deviation magnitude under any similarity-invariant norm* (a norm derived from the dynamics, such as the F -induced metric, rather than a fixed external basis). Each is annihilated by the similarity action.

The architectural content — the basis T , that is, *which* internal coordinates implement the correction — lives entirely in the similarity orbit and is annihilated by the on-policy observation map. $\square$

Scope of “identical.” The claim restricts to *similarity-invariant* summaries of the on-policy observation law. Fixed-external-basis quadratic forms in the raw residual are *not* in general similarity-invariant — for instance,

$$\mathbb{E} \|\delta\|^2 = \text{tr } \Pi, \quad \text{tr}(T \Pi T^\top) \neq \text{tr } \Pi \text{ in general} \quad (\text{AL.4})$$

(equality only when T is orthogonal). The on-policy observer who consumes (α, R) summaries and the innovation spectrum sees only similarity-invariant content; an observer with access to a fixed external basis can in principle break the orbit by observing the raw covariance, but that is a *different* observation regime than the on-policy summary access the no-go is stated under.

The General Sub-Scope (robust qualitative).

Outside the linear-Gaussian sub-scope, the no-go is robust-qualitative rather than exact. Model the agent as a structural causal model over its internal state space (per #scope-agent-identity’s token-level commitment to singular-trajectory agents). The Bareinboim, Correa, Ibeling & Icard (2022) Causal Hierarchy Theorem then applies at the agent-as-SCM layer: two SCMs over the agent’s state space agreeing on Level-1 observation data — the on-policy summary distribution — cannot in general be distinguished on Level-2 questions about $d\mathcal{O}(\cdot)$ on the agent’s internal coordinates.

Derived (architecture-noidentifiability-general, from causal hierarchy theorem applied at agent-as-SCM layer), robust qualitative.

The dual-anchor structure (sharp Kalman-Ho for the linear-Gaussian sub-scope; CHT-at-agent-as-SCM for the general case) is the same shape Instance 1 carries — explicit construction sharp / CHT general — and the tier separation is load-bearing: collapsing to “exact” overclaims; collapsing to “robust-qualitative” understates the sub-scope’s sharpness.

Mechanism: Instance-2 Fisher Null on a Lie-Group Fiber.

The no-go’s *generating* group action is state-space similarity $F \mapsto T F T^{-1}$ on the realization manifold. Sylvester’s law of inertia (5.1 §Discussion) governs metric *congruence* $\mathcal{G} \mapsto S^\top \mathcal{G} S$, a different group action; the no-go therefore does **not** inherit Sylvester at the generating-action layer.

Derived (mechanism-reduction-to-instance-2, from Fisher-information factorization through similarity-invariants).

The reduction to Instance-2 mechanism proceeds via Fisher geometry. Parameterize the agent by

$$\theta = (\text{similarity-invariants, orbit-coordinates}) \quad (\text{AL.5})$$

— formally, the realization manifold fibered over the transfer-function manifold by the $GL(n)$ similarity group. By the §“No-Go Theorem” derivation, the on-policy observation law depends on θ only through the similarity-invariant base. The Fisher information of the on-policy observation channel, as a quadratic form on the full realization manifold, is therefore *identically zero* on the entire tangent distribution of the similarity fiber. That is the Instance-2 structure (Fisher rank-deficiency along a structurally-forced indeterminacy manifold; R Prop B.7), with the manifold being a $GL(n)$ Lie-group fiber rather than the mixture-indeterminacy manifold of Instance 2.

The agent’s *escape-irreducibility under its metric/coordinate freedom* is then Sylvester at one remove — identically to Instance 2. Once the rank-deficient observed-Fisher is recognized, no reparameterization of the observation model refills the rank-deficient direction; congruence preserves inertia.

The two group actions and their roles, made precise:

- The **similarity action** $T(\cdot)T^{-1}$ on the realization manifold *generates* the indistinguishable architectural pair (A, A') . This is what makes the floor a floor at all, and it is *not* Sylvester — it is the fiber of the realization-to-transfer-function bundle.
- The **congruence action** $S^T(\cdot)S$ on the observed-Fisher governs the agent’s potential escape by reparameterizing its observation model. This *is* Sylvester, and it forbids the escape, identically to Instance 2.

This no-go is therefore a member of the rank-collapse subclass of #disc-identifiability-floor ({Instance 1, Instance 2, Instance 4} via Sylvester at one remove) rather than Instance 3’s projection / Schur-complement obstruction. The additional structure beyond Instance 2 is that the rank-deficient direction is a Lie-group fiber.

Boundary Characterization: Three Structurally Distinct Escapes.

The no-go’s scope is (i) on-policy, (ii) in-regime, (iii) summary-only. Each scope-condition violation corresponds to an AAT capability that admits identification:

Derived (escape-routes, from no-go scope conditions).

Table AL.1

<i>Route</i>	<i>Scope violated</i>	<i>AAT capability</i>	<i>Detection strength</i>
(a) Loop-interventional access	(i) on-policy	7.3 — $do(\cdot)$ on the agent’s input	Sharp in linear-Gaussian sub-scope; robust qualitative general
(b) Higher-moment / out-of-regime	(ii) in-regime + (iii) summary-only	Beyond sub-scope α : higher-moment statistics of the innovation; out-of-regime sampling	Provably void in linear-Gaussian sub-scope; informative in nonlinear sub-scope β
(c) Architecture instrumentation	(iii) summary-only (white-box)	Direct read of the update rule or internal state; same kind as Instance 2’s “instrument the latent”	Strongest when available

Route (a) — loop-interventional access (sharp). A $do(\cdot)$ intervention on the agent’s input perturbs the system along *similarity-fiber directions that on-policy data never excites*. The interventional response (correction-function Jacobian probed off the on-policy manifold) is *not* similarity-invariant: TFT^{-1} and F respond differently to a probe in a fixed external basis. The intervention generates Level-2 data; the §“No-Go Theorem” derivation is a Level-1-only statement. Genuinely escapes the orbit degeneracy.

Route (b) — higher-moment / out-of-regime (sub-scope-bounded). In the nonlinear sub-scope β , two agents matched at (α, R) and second-order generally differ at moments of order ≥ 3 . In the linear-Gaussian sub-scope this route is *provably void*: a Gaussian innovation carries no information beyond second order, so higher-moment observation adds nothing — the similarity orbit remains invisible. The “escape exists here, provably not there” boundary is exactly the kind of sharp scope-statement that makes a floor instance informative.

Route (c) — architecture instrumentation. Direct read of the update rule or internal state breaks the black-box scope: the observer now has access to the basis T itself, not just its annihilated projection onto the on-policy observation law. Structurally the same kind as Instance 2’s escape (i) (“instrument the latent” — R Prop B.7); genuinely distinct from (a) and (b) in *what the observer must have* (white-box access versus interventional access versus moment-level access).

Route (d), horizon extension under the same policy — does not escape. Passive observation at longer horizons under the *same* on-policy execution generates more *samples* of the same observation law; the §”No-Go Theorem” derivation is a statement about the *law* (the architectural content is annihilated by the spectrum map), not about the sample size. Horizon extension that takes the agent *out of the current regime* is no longer the same policy in the same regime — that is a regime change, which is structurally an exogenous intervention-like event, collapsing into route (a)/(b). Route (d) is therefore either void (same regime) or a special case of (a)/(b); the count of structurally distinct escapes is **three**, not four.

Why Fano Is the Finite-Sample Refinement, Not the Exact Anchor.

A Fano-style anchor (Fano 1961, *Transmission of Information*) gives a lower bound on prediction error probability *given* a bound on the channel mutual information $I(A; \text{obs})$ between the architectural variable and the observation. Under the similarity-orbit construction, the on-policy observation law is *identical* for A and A' — not merely close — so $I(A; \text{obs}) = 0$ exactly. Fano degenerates to the trivial bound “error probability $\geq$ prior,” vacuous and strictly weaker than the *exact* Kalman-Ho indistinguishability statement.

Derived (fano-degeneracy-at-zero-information), exact.

Fano is therefore the right tool for the *approximate / finite-sample* version of the no-go — architectures that are close but not equal in innovation spectrum, $I > 0$ small — and the wrong tool for the *exact* population-level no-go. The exact anchor is Kalman-Ho; Fano is honest open follow-on work on the finite-sample refinement, distinct from the floor itself. (Cross-reference: #der-agent-opacity Working Notes record the parallel finding from the observer-side H_b task.)

Strengthened Consequence.

The no-go elevates #der-loop-interventional-access to a load-bearing role at the *agent-internal* layer — the third semantically distinct deployment mode after:

Derived (third-mode-of-interventional-access).

1. **Mode 1: agent-self-intervention** — Instance 1’s escape, breaking a *causal-sufficiency* degeneracy among the agent’s action propositions (9.2).
2. **Mode 2: observer-on-sub-agent intervention** — Instance 3’s escape, breaking a *coupling-sign* degeneracy across composite sub-agents (J).
3. **Mode 3: observer-on-agent-input intervention** — this segment’s escape, breaking a *similarity-orbit* degeneracy on the agent’s internal architecture.

The three modes are *different group-theoretic objects*: each breaks a Level-1 degeneracy of a distinct kind, yet shares the load-bearing content “Level-2 data is the unique broadly-available violation of the Level-1 no-go.” This converts the Instance-4 floor from “yet another no-go” into structural confirmation that #der-loop-interventional-access carries the same load-bearing role at *three* distinct AAT layers (action, composite, architecture), each escape preserving the framework’s existing scope honesty rather than introducing new machinery.

Tier. *Exact* in the linear-Gaussian / Kalman sub-scope (α_1), anchored by Kalman-Ho canonical-form non-uniqueness. *Robust qualitative* in the general sub-scope, anchored by CHT-at-agent-as-SCM. **The exact / robust-qualitative tier boundary is load-bearing and must be carried consistently** — collapsing to flat “exact” overclaims the general case; collapsing to flat “robust qualitative” understates the sub-scope’s sharpness. The mechanism reduction (§”Mechanism”) is exact: the Fisher-null factorization through similarity-invariants follows in elementary steps from the §”No-Go Theorem” derivation.

The escape-count statement (“three structurally distinct escapes, with the proposed fourth collapsing into the first”) is *exact* — the collapse of route (d) into (a) and the void of route (b) inside the linear-Gaussian sub-scope are both elementary consequences of the no-go statement itself. The escape-machinery mappings are *robust qualitative*, inheriting the tier of the cited segments (7.3, 16.2).

Epistemic Status.

Conditional. Max attainable: *exact* for the Kalman-Ho derivation in the linear-Gaussian sub-scope; *robust qualitative* for the general case via CHT-at-agent-as-SCM; *exact* for the mechanism reduction to Instance-2 Fisher-null on a Lie-group fiber; *exact* for the escape count (three, not four) and for the Fano-degeneracy-at- $I = 0$ observation; *robust qualitative* for the escape-machinery mappings.

Load-bearing: - The Kalman-Ho construction makes the architectural-identifiability no-go *exact in a precisely characterized sub-scope* — replacing the prior sketch-grade framing (16.2 Working Notes’ candidate-status for a generic observer-side floor). - The mechanism reduction to Instance-2 Fisher-null on a $GL(n)$ fiber places the no-go in the rank-collapse subclass {Instance 1, Instance 2, Instance 4} via Sylvester at one remove, repairing #disc-identifiability-floor’s Sylvester-mechanism taxonomy to three rank-collapse members plus one projection/closure member ({Instance 3}). - The three-mode classification of #der-loop-interventional-access (causal-sufficiency / coupling-sign / similarity-orbit) is a structural unification across three identifiability-floor instances.

Not established: - The general (non-linear-Gaussian) sub-scope at *exact* tier. CHT-at-agent-as-SCM is robust-qualitative; an explicit construction-grade no-go in general sub-scope β is honestly open. - The finite-sample Fano refinement (architectures close-but-not-equal in innovation spectrum, $I(A; \text{obs}) > 0$ small). Open follow-on direction, distinct from the population-level no-go derived here. - Sharp characterization of the boundary between similarity orbits and behavioral equivalence classes when the agent’s nonlinear dynamics admit *partial* spectrum identifiability (the sub-scope- β moment-cascade question).

Honest Limits.

- **Linearization scope.** The Kalman-Ho construction operates on the *linearized* closed-loop residual dynamics. Sector-condition-satisfying agents (A) admit such linearization in a basin around the singular trajectory; the no-go’s exact tier inherits that scope. Outside the basin, the general-case CHT anchor applies but the construction is not closed-form.
- **Action-marginal invariance vs. internal-coordinate freedom.** Per #scope-agent-identity, the action space is coordinate-free and observed quantities are parameterization-invariant. The similarity orbit is a *different* invariance — internal to the realization manifold, distinct from action-marginal parameterization-invariance — and the no-go is about identifying the internal architecture, not about action-level identifiability.
- **Closed-loop residual dynamics, not open-loop plant.** The no-go is stated for the closed-loop δ -dynamics (3.4) — what the on-policy observer actually sees. Open-loop plant identification from interventional experiments is route (a) and is *not* subject to this no-go.

Discussion.

Position in the identifiability-floor pattern. This no-go is the fourth instance of #disc-identifiability-floor and the third member of the rank-collapse subclass via Sylvester at one remove. The Sylvester-mechanism Discussion of the meta-segment (5.1 §Discussion) is repaired by this addition: the three-rank-collapse taxonomy is {Instance 1, Instance 2, Instance 4} unified by congruence-preserves-inertia at the escape-irreducibility layer; Instance 3 stands alone as the projection/closure obstruction. The plurality-of-mechanism that makes the floor pattern a *family* rather than a single theorem is preserved (Sylvester for rank-collapse, Schur/projection for composition); the Sylvester-side is strengthened from two to three members.

Why Instance-2's mechanism, not new. The Fisher-null direction here is a $GL(n)$ Lie-group fiber rather than the single mixture-indeterminacy manifold of Instance 2, but the *kind* of obstruction is identical: a rank-deficient observed-Fisher along a structurally-forced indeterminacy manifold. Naming it a new mechanism would be inflation — the AAT-distinctive recognition is precisely that the architectural d.o.f. AAT cares about (#der-agent-opacity's "which internal mechanism") is the similarity orbit and that the orbit is annihilated by the on-policy observation map; the obstruction theorem itself is the same shape as Instance 2. Naming it Sylvester-for-free would be the opposite inflation — similarity and congruence are different group actions. The honest position is "Instance-2 mechanism on a Lie-group fiber, Sylvester at one remove on the escape."

Relationship to #der-agent-opacity's observer-indexed H_b . This segment's no-go is structurally distinct from a generic observer-prediction-of-action floor: the architectural d.o.f. is invisible *at the observation-law level* regardless of observer instrumentation, not because of a particular observer's filtration. An observer with arbitrarily fine $\mathcal{F}_B^t$ over the on-policy summary distribution still cannot escape the similarity orbit; the observer-indexing of H_b is orthogonal. The natural relationship is at the meta-pattern level (both name floors in the family of #disc-identifiability-floor); the *quantities* are distinct (H_b is observer-conditional action entropy; this segment's invariant is the similarity orbit).

Conservative integration with the Sylvester recognition. The 2026-05-14 Sylvester-recognition Finding in #disc-identifiability-floor taxonomized two rank-collapse floors (Instances 1, 2) and explicitly said "*the floors do not share one mechanism — Sylvester for rank-collapse, a projection/closure obstruction for composition.*" This segment's addition does not contradict that recognition; it extends the rank-collapse subclass to three members and preserves the projection/closure distinction for Instance 3. The plurality-as-feature framing of the meta-segment's Discussion is intact.

Multi-agent / Hafez bridge. The architectural d.o.f. that this no-go identifies as un-recoverable from on-policy summary data is structurally adjacent to #der-agent-opacity's H_b in this sense: both quantities concern the *world's view of the agent's internal architecture* through information-theoretic limits. The Hafez et al. 2026 IDT pattern (Information Digital Twin, bi-predictability P) — operationally a low- H_b -preserving observation channel — is a structurally-instrumented escape, and its 89% perturbation-detection performance is consistent with this segment's route (c) "architecture instrumentation" classification.

Findings.

Architecture No-Identifiability from On-Policy Summary Data — the Fourth Identifiability Floor. **Brief:** Imagine two agents that both keep a desk tidy by the same rules: same overall reaction speed, same tolerance for clutter, same long-run mess level. One agent reasons by tracking pile heights; the other tracks distances between piles. Watching them work — under their normal habits, summary statistics only — you literally *cannot tell which one is running*. The visible behavior is the same, the variability is the same, the rate of clean-up is the same, the cluttered ceiling is the same. Different inner mechanisms, identical outer signature. There are only three honest ways to tell the agents apart: poke one in a way it wouldn't normally do (intervention), look at fine-grained details its rules would not normally generate (out-of-regime / higher-moment observation), or open the hood and read the wiring (white-box access). And a fourth pretender — "just watch longer" — does *not* work, because watching longer accumulates more samples of the same signature, never the missing distinction. This is the floor: it is structurally impossible to recover the inner-mechanism

information from the outer summary alone, and the only paths forward are exactly the three named escapes, each of them violating a clearly-stated scope condition.

Impact: Closes the long-contested fourth instance of #disc-identifiability-floor with a sharp, dual-anchored no-go (Kalman-Ho linear-Gaussian sub-scope, exact; CHT-at-agent-as-SCM general, robust qualitative). Extends the Sylvester-mechanism rank-collapse subclass to three members {Instance 1, Instance 2, Instance 4} via the at-one-remove escape-irreducibility, while preserving the projection/closure distinctness of Instance 3 — the plurality-of-mechanism that makes the floor pattern a presentational family rather than a single theorem stays intact, sharpened. Elevates #der-loop-interventional-access to a third semantically-distinct deployment mode (agent-internal architecture, alongside agent-self-intervention and observer-on-sub-agent), unifying the strengthened-consequence content of three identifiability-floor instances under one piece of machinery. The result mathematically settles the Regime-C confound of the rho-decomposition ($3.5 + \mathbb{W}$) as the same object projected onto the disturbance-statistic coordinate — one decision, not two.

Novelty Claim: *Claim derivation* of the architecture-noidentifiability floor with named scope (linear-Gaussian sub-scope, on-policy summary access) and named dual anchor (Kalman 1963 / Ho-Kalman 1966 / Anderson-Moore 1979 §10.4 for the sharp sub-scope; Bareinboim, Correa, Ibeling & Icard 2022 for the general case). The Kalman-Ho canonical-form non-uniqueness theorem is classical (1963/1966); the AAT-distinctive content is *recognizing* that the architectural degree of freedom the framework cares about — #der-agent-opacity’s “which internal mechanism” — *is* the similarity orbit, and that the orbit is annihilated by the on-policy observation map. *Claim differentiation* on mechanism: the obstruction is Instance-2’s Fisher-null along a structurally-forced indeterminacy manifold (the manifold being a $GL(n)$ Lie-group fiber) rather than a new mechanism, with Sylvester preserving inertia on the escape side identically to Instance 2 — a *member* of the rank-collapse subclass, not a new mechanism family. *Claim recognition* that the previously-proposed Fano anchor degenerates at $I(A; \text{obs}) = 0$ (vacuous lower bound on prediction error under a zero-information channel), so Kalman-Ho is the exact-population anchor and Fano is the finite-sample refinement — a demonstrated no-go on the proposed Fano anchoring, not an assertion.

Related Work:

Table AL.2

<i>ASF concern</i>	<i>Prior-art language</i>	<i>Relationship / Positioning</i>
State-space realization non-uniqueness up to similarity (the sharp sub-scope anchor)	Kalman, R. E. (1963), “Mathematical description of linear dynamical systems,” <i>J. SLAM Control</i> 1(2):152–192; Ho, B. L. & Kalman, R. E. (1966), “Effective construction of linear state-variable models from input/output functions,” <i>Regelungstechnik</i> 14:545–548; Anderson, B. D. O. & Moore, J. B. (1979), <i>Optimal Filtering</i> , Prentice-Hall, §10.4 — minimal realizations sharing an innovation spectrum form a similarity-equivalence class	<i>formal antecedent</i> — adopted with citation; AAT-distinctive recognition is that the architectural d.o.f. of #der-agent-opacity <i>is</i> the similarity orbit
Causal Hierarchy Theorem (the general-case anchor)	Bareinboim, E., Correa, J. D., Ibeling, D. & Icard, T. (2022), “On Pearl’s Hierarchy and the Foundations of Causal Inference,” in <i>Probabilistic and Causal Inference: The Works of Judea Pearl</i> , ACM	<i>formal antecedent</i> — applied at the agent-as-SCM layer; same anchor Instance 1 uses, here in the agent-internal-architecture setting
Fano’s inequality (finite-sample refinement, not the exact anchor)	Fano, R. M. (1961), <i>Transmission of Information</i> , MIT Press; Cover, T. M. & Thomas, J. A. (2006), <i>Elements of Information Theory</i> , Wiley, §2.10	<i>adjacent</i> — vacuous at $I = 0$ for the exact construction; the right tool for the finite-sample architectures-close-but-not-equal version of the no-go (honest open)

<i>ASF concern</i>	<i>Prior-art language</i>	<i>Relationship / Positioning</i>
Sylvester’s law of inertia (escape-irreducibility, at one remove)	Sylvester, J. J. (1852), <i>Phil. Mag.</i> 4(23):138–142; Horn, R. A. & Johnson, C. R. (2013), <i>Matrix Analysis</i> (2nd ed.), Thm 4.5.8	<i>formal antecedent</i> — applied here to the <i>escape</i> side (congruence on the observed-Fisher), distinct from the generating action (similarity on the realization manifold). Sylvester forbids the escape by reparameterization identically to Instance 2
Fisher-information reparameterization law	Lehmann, E. L. & Casella, G. (1998), <i>Theory of Point Estimation</i> (2nd ed.), §2.5 — $\mathcal{G}_\varphi = S^T \mathcal{G}_\theta S$ as the standard congruence law	<i>formal antecedent</i> — the law that makes the escape-side reparameterization a congruence action
Hafez et al. 2026 / IDT bi-predictability	Hafez, A. M., Khan, F. M. & Iqbal, M. (2026), <i>A Mathematical Theory of Agency and Intelligence</i> — IDT pattern, bi-predictability P	<i>adjacent</i> — IDT-as-architecture-instrumentation is route (c) of this segment’s escape structure; the 89% IDT detection rate is consistent with white-box access escaping the orbit

Search Log: - 2026-05 (*targeted*): Kalman 1963 / Ho-Kalman 1966 / Anderson-Moore 1979 §10.4 verified as the canonical-form-non-uniqueness anchors; CHT-at-agent-as-SCM application of Bareinboim et al. 2022 verified consistent with Instance 1’s general-sub-scope use. Fano 1961 $\leftrightarrow$ Cover-Thomas §2.10 verified for the finite-sample refinement framing. Sylvester 1852 + Horn-Johnson Thm 4.5.8 + Lehmann-Casella §2.5 inherited from the 2026-05-14 Sylvester-recognition Finding in #disc-identifiability-floor. - 2026-04 / 2026-05 (*intuition-only*) on whether the cross-instance recognition “the architectural d.o.f. AAT cares about *is* the similarity orbit, annihilated by the on-policy observation map” has been articulated in the system-identification or causal-inference literatures as the AAT-internal-agent-theory framing applied to multi-agent identifiability. The constituent pieces (Kalman-Ho canonical-form non-uniqueness; CHT-at-agent-as-SCM; Fisher rank-deficiency along Lie-group fibers) are textbook; the cross-instance unification as the fourth member of an integrated agent-theoretic floor pattern appears to be a presentational contribution under cursory search. Targeted future search candidates: system-identification literature on the agent-theoretic interpretation of canonical-form non-uniqueness; observability / identifiability literature for nonlinear agents; mechanism-grounded identifiability in multi-agent reinforcement learning.

Working Notes.

- **Provenance.** The architectural-noidentifiability floor was previously sketched in spikes/spike-neutral-drift-endogenous-coupling-strengthening-2026-04-24.md §8 / §10.1 (candidate, Fano anchor proposed but uncompleted) and surfaced from the disturbance-statistic coordinate as the Regime-C confound of spikes/.integrated/spike-rho-structure-recheck-2026-05-18.md §3 (same object). The 2026-05-18 resolution spike spikes/.integrated/spike-identifiability-floor-instance4-resolution-2026-05-18.md re-derived the no-go independently and supplied the Kalman-Ho construction, the mechanism reduction to Instance-2-on-a-fiber, the escape audit, and the Fano-degenerates-at- $I = 0$ no-go on the proposed Fano anchor; the math-gate trail at spikes/.routing-trail/SPIKE-VERIFY-471802/ (confirmer $\neq$ author) cleared the substantive math conditional on three named repairs that are applied in this segment: (i) the “identical” claim restricted to similarity-invariant summaries (not raw moments); (ii) the Lyapunov-equation sign matching #deriv-sector-condition’s standard form ($F\Pi + \Pi F^T = \sigma_w \sigma_w^T$); (iii) the exact / robust-qualitative tier boundary carried consistently rather than collapsed to flat exact.
- **Finite-sample Fano refinement is honestly open.** The Fano-style bound on the *approximate* version (architectures close-but-not-equal in innovation spectrum, $I(A; \text{obs}) > 0$ small) is real research, distinct from the population-level Kalman-Ho no-go. Candidate spike for a future cycle; not blocked by anything here.
- **General non-Gaussian sub-scope at exact tier.** The CHT-at-agent-as-SCM general-case anchor is robust-qualitative. An explicit construction-grade no-go in the nonlinear sub-scope β would upgrade the general case to exact-with-conditions; honestly open.
- **The mechanism reduction’s relationship to operator-sector machinery.** The §”Mechanism” section’s recognition that the Fisher-null here is a rank-deficiency on a $GL(n)$ Lie-group fiber places it in the stability-certificate spine’s *boundary facet*

(5, 5.1) at the same level as Instance 2's mixture-indeterminacy manifold. Whether the Lie-group-fiber generalization of Instance 2's mechanism is worth its own meta-statement (a "Fisher-null-along-structurally-forced-manifold" sub-pattern of the rank-collapse subclass) is open architectural taxonomy.

Appendix AM

Derivation: Fisher-Whitened Edge Update Under Correlated Evidence

Under L1'/L2 correlated-evidence regimes, the default log-odds edge-update (from #deriv-edge-update-natural-parameter) retains correct *direction* — the angle between log-odds gradient and natural gradient never exceeds 45° at finite correlation ρ , so B1 directional fidelity (4.3) is never actively violated — but its *magnitude alignment* degrades by a factor $\sqrt{1 - r^2}$ in the sector constant. The Fisher-whitened correction restores sharp B1 on the Fisher-weighted inner product. Under the (PI) parameterization-invariance axiom named in #scope-agent-identity and promoted to a primary instance of #disc-additive-coordinate-forcing via Čencov 1982, Fisher whitening is **AAT-internally derivable** rather than externally imported. The result adds sub-scope α_3 (correlated evidence + Fisher-whitened update + Bayesian coherence $\rightarrow$ A2' derived) to the A2' partition and composes cleanly with the meta-gain machinery of #deriv-adaptive-gain-dynamics (Fisher whitening is a special case of meta-gain with $K_t = \mathbf{I}^{-1}(\lambda_t)$).

The framework's posture splits cleanly. *Without* the (PI) commitment, the Fisher-whitened correction is an externally imported choice from information geometry. *With* (PI), Čencov 1982's uniqueness theorem *forces* the Fisher metric as the unique Markov-invariant metric on the statistical manifold, and the whitened-update form is then the AAT-internally canonical correction for directional-fidelity preservation across reparameterizations. Operational prescription: where correlated evidence is the dominant regime, Fisher whitening recovers sharp persistence guarantees at the cost of meta-gain machinery; where it is mild or rare, the default log-odds update with $\sqrt{1 - r^2}$ magnitude-alignment degradation suffices. Sibling at the model-parameter-update layer is AR (update *magnitude* under Fisher-local regime); together they make the Fisher-local invariance regime AAT-internally complete.

Formal Expression.

Angle characterization under L1' correlation.

For a block-correlated two-edge evidential model (soft-facilitator Prop B.7 parameterization), the Fisher information matrix has off-diagonal entry

Derived (angle-bound-finite-correlation).

$$r = \frac{\theta_C(1 - \theta_C)\Delta_1\Delta_2}{\sigma_1\sigma_2} \in [-1, 1], \quad (\text{AM.1})$$

where θ_C is the latent-cause probability, $\Delta_j = p_{j|C} - p_{j|\neg C}$ is the separability gap, and σ_j is the marginal-edge standard deviation. The Euclidean angle between the log-odds update direction and the natural-gradient direction (at single-edge plan residual) satisfies

$$\theta_{\text{LO-NG}} = \arccos \frac{1}{\sqrt{1+r^2}} \leq 45^\circ \quad \text{for all } r \in [-1, 1]. \quad (\text{AM.2})$$

The log-odds update **never flips sign** under any finite correlation — B1 is never actively violated, only degraded in magnitude. Under the operational claim that “the default signal function needs validation under correlated failures,” the finding is: **the default direction is preserved; only magnitude alignment degrades**. The sector constant in B1 degrades by a factor $\sqrt{1-r^2}$ under Euclidean B1; the Fisher-weighted B1 recovers the unperturbed sector constant.

Fisher-whitened update rule.

In log-odds coordinates $\lambda \in \mathbb{R}^{|E|}$ with Fisher information $\mathbf{I}(\lambda)$, the Fisher-whitened edge update is

Definition (Fisher-whitened-update).

$$T_{\text{FW}}(\lambda) = \lambda - \eta_{\text{edge}} \cdot \mathbf{I}^{-1}(\lambda) \cdot \mathbf{J} \cdot (\hat{P}_\Sigma(\sigma(\lambda)) - y_G) \quad (\text{AM.3})$$

(compared to the Euclidean log-odds update $T_{\text{edge}}(\lambda) = \lambda + \eta_{\text{edge}} \cdot \text{diag}(t) \cdot \mathbf{J}(y_G - \hat{P}_\Sigma)/\|\mathbf{J}\|^2$). The Fisher-weighted inner product $\langle a, b \rangle_{\mathbf{I}} = a^T \mathbf{I}^{-1} b$ makes the update’s directional fidelity invariant under reparameterization of the natural-parameter coordinate.

Two AAT-internal axiom paths.

Path A (B1-parameterization-invariance). Require B1 directional fidelity to be *parameterization-invariant* in the sense of #scope-agent-identity’s (PI) axiom: the theorems about sub-scope α derivation in #der-gain-sector-bridge should not depend on arbitrary coordinate choices for M_t ’s natural parameters. Under (PI), B1 sub-scope α partition is coordinate-invariant iff the inner product defining the directional-fidelity condition is the Fisher metric (Čencov 1982 uniqueness theorem under (PI); extended by Ay-Jost-Lê-Schwachhöfer 2017). The Fisher-weighted inner product is therefore *forced* by (PI), and the Fisher-whitened update is the AAT-internally derived correction direction for directional-fidelity preservation across parameterizations.

Derived (Fisher-whitening-from-B1-parameterization-invariance).

Path B (Lyapunov-coordinate-matching via adjacent-family classification). In the adjacent-family framing of #disc-additive-coordinate-forcing, the Lyapunov coordinate is *matched* (not forced) to the sector condition. For natural-gradient updates, the canonical Lyapunov is Fisher-weighted (Amari 1998, “Natural gradient works efficiently in learning,” *Neural Computation* 10); this matches the geometry of the update operator. The two paths converge on the same Fisher-weighted result; Path A forces it via axiomatics, Path B confirms it via adjacent-family coordinate-matching.

Under L0 (no correlation), $r = 0$ and $\mathbf{I}$ is diagonal — Fisher whitening is vacuous (reduces to the existing Euclidean log-odds update). The axioms pick out Fisher whitening *uniquely* only under L1’/L2 (correlated evidence) regimes; they are vacuously satisfied under L0.

New sub-scope α_3 .

Correlated-evidence + Fisher-whitened update + Bayesian coherence yields A2' *derived*:

Derived (sub-scope-alpha-3).

$$(T_{\text{FW}}(\lambda) - \lambda^*)^T \mathbf{I}(\lambda^*)^{-1} (\lambda - \lambda^*) \geq \beta \|\lambda - \lambda^*\|_1^2 \quad \text{with } \beta = \eta_{\text{edge}} \mu \quad (\text{AM.4})$$

on the Fisher-metric sector region around the fixed point $\lambda^* = \text{logit}(\theta^*)$. The directional-fidelity proof of #der-gain-sector-bridge carries over with $\kappa(\mathbf{I})$ as Euclidean-transfer penalty — structurally identical to the Kalman case's $(P^-)^{-1}$ -norm weighted-norm treatment. The Kalman case in #der-gain-sector-bridge's Verified Instances is a special case of Fisher-weighted sector.

Sub-scope α_3 extends the refinement introduced by #deriv-adaptive-gain-dynamics: - **Sub-scope α_1** : fixed-gain, independent evidence (Euclidean B1 applies; the existing A2' partition). - **Sub-scope α_2** : adaptive-gain, independent evidence (meta-gain conditions (MG-1)–(MG-4) per #deriv-adaptive-gain-dynamics). - **Sub-scope α_3** : fixed-gain, correlated evidence (Fisher-whitened update; this segment).

Additional composition: adaptive-gain + correlated-evidence (sub-scope α_4) would compose #deriv-adaptive-gain-dynamics meta-gain machinery with Fisher whitening at the primary state; open.

Connection to meta-gain as special case.

The Fisher-whitened update is a meta-gain law in the sense of #deriv-adaptive-gain-dynamics with $K_t = \mathbf{I}^{-1}(\lambda_t)$ — a *degenerate* special case where the meta-gain is a deterministic function of the primary state rather than an independently learned variable. All four meta-gain conditions (MG-1)–(MG-4) are satisfied trivially: - (MG-1) Meta-gain sector: $\mathbf{I}^{-1}$ is symmetric-positive-definite on the interior of the natural-parameter domain. - (MG-2) Meta-gain bounded: $\|\mathbf{I}^{-1}\|$ bounded on compact natural-parameter subsets. - (MG-3) Smoothness: $\mathbf{I}^{-1}$ depends smoothly on λ for exponential families. - (MG-4) Primary-meta coupling bounded: $\mathbf{I}^{-1}$'s state-derivative in the drift direction is bounded for exponential families.

Derived (Fisher-whitening-as-special-meta-gain).

This hands #deriv-adaptive-gain-dynamics a **concrete second instance** of derivable meta-gain alongside adaptive-Kalman-with-Mehra-estimator (its primary instance). The machinery of meta-gain composition via #result-sector-persistence-template (augmented-state Lyapunov) applies directly.

L2 regime: candidate third #disc-identifiability-floor instance or strengthening of Instance 2.

Under L2-latent regimes (unobservable correlation structure beyond what Fisher information can resolve), Fisher whitening fails: the correlation sub-block of Fisher is rank-deficient (Cramér-Rao floor on unobservable parameters). This parallels #disc-identifiability-floor Instance 2's L1'-unobservable-C Fisher-rank-1 obstruction — potentially a new L2 instance, or potentially a generalization of Instance 2 to higher correlation orders. Open whether these are distinct floors or a single unified obstruction.

Hypothesis (L2-latent-floor).

L2-degenerate (perfect correlation, $r \rightarrow 1$) is a *structural* collapse — it requires DAG repair (edge merging) rather than update repair. This sits outside the Fisher-whitening framework.

Epistemic Status.

Conditional. Max attainable: *exact* under Path A ((PI) + Čencov combined with B1 directional fidelity) + L1'/L2 regimes + Bayesian coherence. The derivation is textbook information-geometry (Amari 1998; Amari-Nagaoka 2000) plus the AAT-internal (PI) axiom from #scope-agent-identity. What is AAT-distinctive is (a) the identification of (PI)/Čencov as the forcing machinery rather than merely a preferred choice; (b) the sub-scope α_3 labeling within the A2' partition; (c) the composition with #deriv-adaptive-gain-dynamics' meta-gain framework as a degenerate special case.

The angle bound is $\leq 45^\circ$ for all finite r . This is a sharp load-bearing observation: AAT's existing default log-odds signal function is *directionally robust* under correlated failures. The bias formula in #deriv-l1-update-bias gives the quantitative magnitude-bias; this segment gives the qualitative direction-preservation. Together they characterize the default signal function's behavior under L1' fully: **direction preserved, magnitude biased, observable-C restores exactness (Prop B.7), unobservable-C floor per F13.**

Without (PI) adoption. If (PI) is not adopted as an AAT axiom, Fisher whitening remains a theorem-imported technique from information geometry. The sub-scope α_3 labeling still applies (correlated evidence + Fisher-whitened + Bayesian coherence $\rightarrow$ A2' derived) but without the AAT-internal forcing. This is the adjacent-family status.

Discussion.

Relationship to the default signal function. #disc-credit-assignment-boundary's default log-odds signal was landed in the 2026-04-22/23 strengthening cycle via the Cauchy-FE uniqueness theorem in #deriv-edge-update-natural-parameter. Under L0 evidence independence, the default signal is exact. Under L1' correlated evidence, the default signal's direction is preserved ($\leq 45^\circ$ angle to natural-gradient direction) but magnitude degrades by $\sqrt{1-r^2}$. This segment's Fisher-whitened update is the correction under (PI)/Čencov; it recovers sharp B1 under L1' by lifting the inner product from Euclidean to Fisher.

Relationship to #disc-identifiability-floor Instance 2. Instance 2 (L1' unobservable-C Cramér-Rao floor) refutes single-channel identification of mixture parameters. Under observable C , Prop B.7 gives the L1' transfer; under unobservable C , no update — including Fisher-whitened — recovers identification. Fisher whitening operates *downstream* of identification: it corrects the update direction given that the Fisher information is known; it does not manufacture Fisher information that the data cannot yield.

Composition with the (PI)/Čencov 4th primary instance. The Fisher-metric cases in #der-gain-sector-bridge's Verified Instances table (Matrix Kalman, Exponential family in natural parameters) are also AAT-internally forced under (PI)/Čencov. Fisher whitening at the edge-update layer of Part II therefore shares the same AAT-internal axiom as the Kalman / exp-family cases in Part I's update machinery. This is a cross-layer consistency: the (PI) axiom operates uniformly across Parts I and II whenever the natural parameters are parameters on a statistical manifold.

Working Notes.

- **Sub-scope α_4 composition** (adaptive-gain + correlated-evidence). Composing #deriv-adaptive-gain-dynamics meta-gain machinery with Fisher whitening gives an augmented-state contraction argument where both the primary state (log-odds) and the meta-state (Fisher information estimate) are updated. Open derivation.
- **L2 regime.** Fisher whitening at L2 (non-tree correlation structure) requires computing higher-order Fisher correlations. For sparse DAG structures with limited correlation order, Fisher-tree approximations may suffice; for dense correlation structures, the computational cost scales exponentially. Scope boundary for Fisher whitening as a practical technique.
- **Relationship to natural-gradient variational inference.** Khan & Lin 2017 (*Conjugate-computation variational inference*) derives natural-gradient VI as a special case of conjugate Bayesian computation. Under (PI)/Čencov, their result becomes

AAT-internally derived within sub-scope α_3 . This composes with `#deriv-variational-sector-condition` if that segment lands — natural-gradient VI would be the $\varepsilon = 0$ limit of ε -fidelity B1.

- **Landing context.** This segment lands from `spikes/.integrated/spike-fisher-whitened-update.md` (2026-04-23 Gap A/B cycle). The AAT-internal axiom path via (PI)/Čencov was validated by the jacobian-b1 strengthening (2026-04-23 follow-up; absorbed in the spike-routing cycle, CHANGELOG 2026-05-17). The broader uniqueness-theorem-machinery discussion clearing the `#disc-additive-coordinate-forcing` discipline lives in that meta-segment.
 - **Sibling segment at the model-parameter layer.** `#deriv-fisher-local-update-gain` derives update *magnitude* under the Fisher-local invariance regime, sibling to this segment's derivation of update *direction* under correlated evidence. Both share (PI)/Čencov via Path A's axiomatic chain. Together they make the Fisher-local invariance regime AAT-internally complete: direction at the edge-update layer and magnitude at the model-parameter-update layer.
-

Appendix AN

Derivation: Regime-Marginal Indistinguishability — Cross-Regime Witness for #disc-identifiability-floor Instance 3 (Broadened Scope)

This segment supplies the parameter-by-parameter witness construction for #disc-identifiability-floor Instance 3's broadened-scope claim (2026-05-21): from sub-agent marginal data alone, the coupling-topology bit distinguishing among R0 / R1 / R2 dynamic regimes (per #disc-dynamic-regime-axis) is unidentifiable. The witnesses are exact for symmetric-matched-Tier-1-scalar sub-agents under standard coupling models (Model D for disturbance; the same models #deriv-critical-mass-composition uses, where the Instance-3 coupling-sign-bit witness is established). The mechanism is rank-collapse / Sylvester-at-one-remove on the marginal-Fisher operator for the topology coordinate, parallel to the Instance-3 coupling-sign-bit mechanism but acting on a higher-dimensional topology coordinate.

The relationship to #disc-identifiability-floor Instance 3 is supporting-derivation — Instance 3 carries the no-go statement, the escape menu, and the strengthened consequence; this segment carries the explicit construction backing the broadened no-go.

Formal Expression.

Setup. Consider $N = 2$ scalar sub-agents indexed $i \in \{1, 2\}$, each with state $q_i \in \mathbb{R}$, base correction rate $\alpha > 0$, observation noise $w_i \sim \mathcal{N}(0, \sigma^2)$, base disturbance bound ρ , tempo $\mathcal{T}$. The composite state is $X^c = (q_1, q_2) \in \mathbb{R}^2$. An external observer B accesses *only* per-sub-agent marginal data: for each i , the marginal stationary distribution $\mu_\infty^{(i)}$, the local mismatch process, and the per-sub-agent update rule $f_M^{(i)}$ as a structural object (white-box on each sub-agent, observed in isolation). The observer does *not* see: the composite-level coupling topology (sign + structure of cross-agent influence), the joint trajectory $\{X_t^c\}$ (only the projection onto per-sub-agent marginals), the game-structure object $\{O^{(j)}\}_{j \neq i}$ from sub-agent i 's frame, or composite-level convergence-rate-class diagnostics.

The inferential question: from per-sub-agent marginal data alone, can B identify the composite's dynamic regime $\mathcal{R}(X^c) \in \{R0, R1, R2\}$ per #disc-dynamic-regime-axis §B?

Witness 1: R0 vs R1 indistinguishability.

Composite A — R0 shared-objective contraction. Sub-agent objective $O_A^{(i)} = -\frac{1}{2}(q_1 + q_2 - T)^2$ for shared target T . Per-agent gradient-flow best-response:

Derived (cross-regime-witness, from Cournot equilibrium structure + shared-target marginal-matching).

$$\dot{q}_i = -\alpha(q_i + q_j - T) + w_i. \quad (\text{AN.1})$$

The composite is R0 per #disc-dynamic-regime-axis §B.0: unique attracting fixed point at $(T/2, T/2)$, joint Lyapunov $V = \frac{1}{2}\|q - (T/2, T/2)\|^2$, exponential convergence at rate α . The stationary marginal distribution of each sub-agent's state, computed from the linearized Ornstein-Uhlenbeck dynamics, is

$$\mu_{\infty,A}^{(i)} = \mathcal{N}(T/2, \sigma^2/(2\alpha)). \quad (\text{AN.2})$$

Composite B — R1 strategic potential-game equilibrium (Cournot). Sub-agent objective $O_B^{(i)} = q_i(a_0 - b(q_1 + q_2) - c)$ — quadratic-profit Cournot in canonical form. The game is a potential game (Monderer-Shapley 1996) with potential Φ_B satisfying $\partial\Phi_B/\partial q_i = \partial O_B^{(i)}/\partial q_i$ for each i . The unique Nash equilibrium $q_{B,\text{Nash}}^*$ is locally attracting under gradient flow, with local-curvature parameter $\lambda_{\min}(\nabla^2\Phi_B(q_{B,\text{Nash}}^*))$ controlling the rate. Linearizing the best-response gradient flow around the Nash and adding the same observation noise, the stationary marginal distribution of each sub-agent's state is

$$\mu_{\infty,B}^{(i)} = \mathcal{N}(q_{B,\text{Nash}}^*, \sigma^2/(2\lambda_{\min}(\nabla^2\Phi_B))). \quad (\text{AN.3})$$

Parameter-matching. Two free parameters on each side: (T, α) for Composite A, (a_0, b, c) for Composite B (with α implicit in the gradient-flow rate). Two matching conditions:

- **Mean-matching:** $T_A/2 = q_{B,\text{Nash}}^*$. Choose $T_A = 2 q_{B,\text{Nash}}^*$.
- **Variance-matching:** $\sigma^2/(2\alpha) = \sigma^2/(2\lambda_{\min}(\nabla^2\Phi_B))$. Choose b in Composite B such that $\lambda_{\min}(\nabla^2\Phi_B) = \alpha$.

Under both conditions: $\mu_{\infty,A}^{(i)} = \mu_{\infty,B}^{(i)} = \mathcal{N}(q^*, \sigma^2/(2\alpha))$ for both i , where $q^* = T_A/2 = q_{B,\text{Nash}}^*$. The marginal *transient* second moments also match under the matched local-curvature condition — both are exponentially-converging Ornstein-Uhlenbeck-type processes per sub-agent in the linearization.

The R0/R1 witness. Composite A satisfies $\mathcal{R}(\Sigma_A) = \text{R0}$ (unique global attractor + global Lyapunov); Composite B satisfies $\mathcal{R}(\Sigma_B) = \text{R1}$ (unique-Nash equilibrium with local Lyapunov-on-deviation, but Nash structure rather than shared-target structure). The marginal-data observer sees identical $\mu^{(i)}$ in both; cannot distinguish $\mathcal{R} = \text{R0}$ from $\mathcal{R} = \text{R1}$. ■

Witness 2: R0 vs R2 indistinguishability.

Composite C — R2 cyclic-distributional regime (matching-pennies under continuous-action smoothing). Take the 2×2 matching-pennies payoff structure, smoothed to continuous actions via softmax-style best-response with temperature parameter τ . The composite has no pure-strategy Nash; the mixed-Nash is at $(1/2, 1/2)$; under no-regret dynamics (Hart-Mas-Colell 2000) the empirical joint distribution converges to the CCE set at rate $O(1/\sqrt{T})$. The per-sub-agent marginal stationary distribution under this regime is a continuous distribution with mean $1/2$ and bounded variance determined by τ and the noise level σ^2 .

Derived (cross-regime-witness, from matching-pennies aggregation + shared-target marginal-matching).

Parameter-matching to Composite A. Choose Composite A's parameters (T_A, α_A) to align with Composite C's per-sub-agent marginal mean ($T_A/2 = 1/2 \Rightarrow T_A = 1$) and variance (choose α_A such that $\sigma^2/(2\alpha_A)$ equals Composite C's marginal variance — solvable for appropriate τ).

The R0/R2 witness. Composite A is R0 (state-variable macro-state, joint Lyapunov). Composite C is R2 (distributional macro-state, no joint Lyapunov on $\mathcal{X}^c$; CCE convergence only). The marginal-data observer sees identical $\mu^{(i)}$ in both; cannot distinguish $\mathcal{R} = \text{R0}$ from $\mathcal{R} = \text{R2}$. This sharpening is qualitatively cleaner than the R0/R1 witness because the *joint* dynamic regimes differ qualitatively — convergent vs cyclic-in-joint-distribution — while the per-sub-agent marginals can be made stationary in both cases by appropriate observation-aggregation: cyclic in the *joint* distribution projects to *stationary* in the marginal distribution when the cycle is over the joint coordinate that the marginal averages out. ■

Mechanism (Sylvester at one remove, same as Instances 1 / 2 / 4).

The composite-level topology coordinate — the bit (σ, p) distinguishing R0 / R1 / R2, where σ is the coupling-sign-bit (cooperative vs adversarial) and p is the potential-existence bit (potential/monotone game-structure vs none) — enters the per-sub-agent marginal observation law $p(q_i \mid \theta_{\text{topology}})$ only through the *fixed-point* and *local-curvature* parameters that the witness constructions above match across composites. Computing the Fisher information of the marginal observation law with respect to θ_{topology} yields a rank-deficient matrix: the topology coordinate has a null direction along the manifold

Discussion.

$$\{\theta : (\text{mean, variance})(\theta) = (q^*, \sigma^2/(2\alpha))\}. \quad (\text{AN.4})$$

By Sylvester's law of inertia (#disc-identifiability-floor Sylvester-recognition Finding), no reparameterization of the marginal-observation coordinate system removes this null direction — the irreducibility is invariant under the observer's representational freedom. The escape requires *rank-augmentation*: observation of a coordinate the marginal projection has annihilated. This is exactly the same mechanism as #disc-identifiability-floor Instances 1, 2, and 4; the present derivation occupies the *same* rank-collapse subclass as Instance 3's original coupling-sign-bit derivation, with the rank-collapse occurring on a topology coordinate of higher dimension than Instance 3's single sign-bit.

The mechanism's *one-remove* status (per #disc-identifiability-floor Instance 4's Mechanism discussion) is preserved: the *generating* action producing the indistinguishable pair is the swap between cooperative-shared-objective coupling and strategic-potential-game coupling (a structural game-theoretic operation), not metric congruence; but the *escape-side* irreducibility under the observer's representational freedom is identical to Instance 2's Fisher-null-on-a-fiber. The architecture of the proof is parallel to Instance 4's, applied to a different topology coordinate.

Epistemic Status.

Conditional on the standard coupling models (Model D for disturbance; the same models #deriv-critical-mass-composition uses for Instance 3's original coupling-sign-bit construction). The two witness constructions are at parameter-by-parameter explicit level for the symmetric-matched-Tier-1-scalar case (R0 vs R1 with $T_A = 2q_{B,\text{Nash}}^*$ and $\lambda_{\min}(\nabla^2\Phi_B) = \alpha$ as the matching conditions; R0 vs R2 with mean-matching and variance-matching against the matching-pennies-with-continuous-smoothing marginal).

Max attainable: *exact* in the symmetric-matched-Tier-1-scalar sub-scope, anchored by the explicit parameter matching above; *robust qualitative* in the general heterogeneous case, inheriting the Liberzon 2003 / Dayawansa-Martin 1999 / Shorten et al. 2007 common-Lyapunov-nonexistence structure that #disc-identifiability-floor Instance 3 (Boundary Characterization) already invokes. The tier statement is the same one Instance 3 carries; the broader topology coordinate does not change the tier.

Discussion.

Why this is a supporting derivation rather than a sixth identifiability-floor instance. The witness constructions in §"Witness 1" and §"Witness 2" *do* derive a cross-regime no-go from per-sub-agent marginal data, but the integrity test on M1 (the five-element distinctness test) finds three of four distinctness criteria failing against the existing Instance 3:

- Same external theorem (Liberzon 2003 / Dayawansa-Martin 1999 / Shorten et al. 2007 — the same anchor Instance 3 invokes; no new external theorem required).
- Same mechanism family (rank-collapse / Sylvester-at-one-remove on the marginal-Fisher operator).
- Four of five escapes inherit verbatim from Instance 3 (a, b, c, d unchanged; only (e) is new).
- Two of three strengthened consequences inherit-with-broadening from Instance 3; only the #deriv-strategic-composition sub-scope α'/β' elevation is genuinely new.

Per the strengthen-first landing posture (~/.claude/memory/epistemic-discipline/strengthen-before-soften.md and its landing-half ~/.claude/memory/epistemic-discipline/integration-is-replacement.md), this segment lands the *math* in canon (math-lives-in-segments discipline) while the *Instance 3 broadening* in #disc-identifiability-floor lands the *no-go statement, escape menu, and strengthened consequence* (integration-is-replacement at the meta-segment scope level). The candidate "Instance 6" framing is deleted; the broadening of Instance 3 carries the strengthened content; this segment carries the explicit witness backing.

The split — math segment plus meta-segment broadening rather than a new instance — preserves M1's instance-count integrity (the instances are *plural in mechanism*, not just plural in surface no-go) while ensuring the derived content lands in canon at full visibility.

Relationship to #disc-dynamic-regime-axis. This segment's witnesses are what concretely back #disc-dynamic-regime-axis's §H "Identifiability floor for regime" — the regime-axis meta-segment names the negative-half complementarity in the abstract; this segment supplies the constructive witness showing the no-go is real and parameter-explicit. The relationship is parallel to how #disc-identifiability-floor Instance 3's coupling-sign-bit witness in #deriv-critical-mass-composition (asymmetric-limit treatment) backs the meta-segment-level no-go statement.

When the dynamic-regime axis advances to meta-segment status, this segment is the canonical backing for the regime-axis floor; the complementarity structure (#disc-separability-pattern ↔ #disc-identifiability-floor parallel for the regime axis) is then:

- *Positive half:* #disc-dynamic-regime-axis four-tier ladder (R0 / R1 / R2 / R3 with their respective Lyapunov machinery).

- *Negative half*: #disc-identifiability-floor Instance 3 (broadened) — backed by this segment’s witnesses.

The double-duty Instance 3 serves both #disc-separability-pattern Contraction ladder *and* #disc-dynamic-regime-axis four-tier ladder; this segment’s witnesses extend that double-duty to the regime-axis level explicitly.

Open extensions.

- **Composite-of-composites case.** The witness constructions handle finite- N matched-symmetric-Tier-1 sub-agents. The composite-of-composites question (regime-identification of meta-composites from sub-composite marginal data) plausibly inherits the Liberzon machinery one level up but is not derived here.
- **Finite-sample refinements.** The §”Witness 1” / §”Witness 2” constructions are exact-population statements (stationary distributions identical). The finite-sample minimax error of regime classification under observation budget T is the Fano-style refinement that the 2026-05-20 spike correctly identified as the *adjacent* tool. Should be derivable from #der-agent-opacity H_b machinery applied at the composite layer; not done here.
- **Mean-field R3 lift.** The §”Witness 1” / §”Witness 2” constructions handle R0 / R1 / R2 finite- N . The R3 mean-field tier (per #disc-dynamic-regime-axis §B.3) has its own identifiability question — from finite- N marginal data, can the limit MFG regime be identified? Plausibly governed by the Lasry-Lions-monotonicity-vs-not bit (the MFG-analog of the potential-existence bit handled here). Tied to the R3 lift in #deriv-strategic-composition Working Notes mean-field extension.

Findings.

Cross-Regime Marginal Indistinguishability (Witness Backing for Broadened Instance 3). **Brief:** When you watch only one player at a time in a multi-agent system, you cannot tell what game they are playing — and “what game they are playing” determines whether the system as a whole will settle to a target, settle to a strategic equilibrium, or cycle in distribution forever. Two concrete worked constructions show this: two firms in Cournot competition that produce identical per-firm production-quantity distributions to a rowing team optimizing for a shared target (different game, same marginal); two players in matching pennies whose mixed strategy produces identical per-player action distributions to the same rowing team (different *kind* of equilibrium, same marginal). The mechanism is the same as for the simpler coupling-sign-bit case: the topology bit lives in a direction the marginal projection annihilates, and Sylvester’s law of inertia says you cannot recover an annihilated direction by changing the units on the ruler — only by looking from a new vantage point (intervening on the topology, or observing the joint trajectory’s convergence-rate class).

Impact: Backs #disc-identifiability-floor Instance 3’s 2026-05-21 broadening from coupling-sign-bit to coupling-topology-bit with explicit parameter-by-parameter witness constructions. Preserves the math-lives-in-segments discipline (the derivation lives in canon, not only in the spike). Provides the canonical worked instantiation for the broadened no-go, parallel to how #deriv-critical-mass-composition’s asymmetric-limit treatment carries Instance 3’s original coupling-sign-bit witness. Operationalizes the “regime is unidentifiable from per-sub-agent marginals” claim that the dynamic-regime axis meta-segment (#disc-dynamic-regime-axis) references as the negative-half of its complementarity structure.

Novelty Claim: *Claim integration* — the witness constructions extend Instance 3’s coupling-sign-bit construction to broader topology coordinates (R0 vs R1; R0 vs R2) using the same Liberzon 2003 / Dayawansa-Martin 1999 / Shorten et al. 2007 anchor and the same Sylvester-at-one-remove mechanism. The Cournot construction (Composite B in §”Witness 1”) is standard Monderer-Shapley 1996 potential-game-as-Cournot; the matching-pennies construction (Composite C in §”Witness 2”) is standard rock-paper-scissors-style cyclic-game smoothing. The contribution is the parameter-matching against R0 shared-target dynamics that produces marginal-indistinguishability across regimes — and the mechanism reduction to Instance 3’s rank-collapse-on-topology-coordinate framework.

Related Work:

Table AN.1

<i>ASF concern</i>	<i>Prior-art language</i>	<i>Relationship / Positioning</i>
Common-Lyapunov nonexistence for switched systems	Liberzon 2003 <i>Switching in Systems and Control</i> §2.1; Dayawansa-Martin 1999 <i>IEEE TAC</i> 44:751; Shorten et al. 2007 <i>SIAM Review</i> 49:545	<i>formal antecedent</i> — same anchor #disc-identifiability-floor Instance 3 invokes; this segment extends the application from coupling-sign-bit to coupling-topology-bit
Cournot duopoly as potential game	Cournot 1838 <i>Recherches</i> ; Monderer-Shapley 1996 <i>Games and Economic Behavior</i> 14:124 §3	<i>formal antecedent</i> — standard potential-game instantiation used as Composite B in Witness 1
Matching pennies / cyclic-game smoothing	Standard textbook (Fudenberg-Tirole 1991 <i>Game Theory</i> §1; Hart-Mas-Colell 2000 <i>Econometrica</i> 68:1127)	<i>formal antecedent</i> — standard cyclic-game template used as Composite C in Witness 2
Sylvester’s law of inertia as escape-irreducibility	Sylvester 1852 <i>Phil. Mag.</i> 4(23):138; cf. #disc-identifiability-floor Sylvester-recognition Finding	<i>formal antecedent</i> — same congruence-invariance machinery the rank-collapse subclass relies on
Multivariate Markov chains with prescribed component behavior but different joint dependence	Bielecki-Jakubowski-Niewkeglowski 2011, <i>Intricacies of dependence between components of multivariate Markov chains: weak Markov consistency and weak Markov copulae</i> (DOI 10.1214/EJP.V18-2238) ¹	<i>closest structural analogue</i> — same-marginals-different-coupling-law construction at the multivariate Markov-chain level; the hidden distinction is filtration-level dependence rather than potential-vs-cyclic game structure, but the formal move (matched per-component behavior, distinct joint law) is the same shape as Witness 1 / Witness 2 at the game-topology level
Marginal-non-determining ergodic measures (stochastic-process side)	Courbage-Hamdan 2001, <i>A family of stationary processes with infinite memory having the same p-marginals</i> (DOI 10.4064/CM90-2-2) ² ; Courbage-Hamdan 1998, <i>An ergodic Markov chain is not determined by its two-dimensional marginal laws</i> (DOI 10.1016/S0167-7152(97)00096-5) ³ ; Hamdan 2000, <i>Markov Chains with Positive Transitions Are Not Determined by Any p-Marginals</i> (DOI 10.1007/S006050070034 — journal-name unverified at landing time) ⁴ ; Hamdan 2002, <i>An ergodic Markov chain is not determined by any p-marginals</i> (DOI 10.1016/S0019-3577(02)80028-3) ⁵	<i>adjacent literature</i> — strong precedent for “finite-order marginals do not determine the generating law” at the stochastic-process level. Different object (Markov chains, not coupled best-response dynamics) but identical formal shape (matched marginals, distinct joint law)

¹The paper and the characterization of its claims are inherited from a second-pass Undermind synthesis (search conducted 2026-05-21); the primary source has *not* been verified at landing time. A primary-source verification spike is queued — see Working Notes above.

²The paper and the characterization of its claims are inherited from a second-pass Undermind synthesis (search conducted 2026-05-21); the primary source has *not* been verified at landing time. A primary-source verification spike is queued — see Working Notes above.

³The paper and the characterization of its claims are inherited from a second-pass Undermind synthesis (search conducted 2026-05-21); the primary source has *not* been verified at landing time. A primary-source verification spike is queued — see Working Notes above.

⁴The paper and the characterization of its claims are inherited from a second-pass Undermind synthesis (search conducted 2026-05-21); the primary source has *not* been verified at landing time. A primary-source verification spike is queued — see Working Notes above.

⁵The paper and the characterization of its claims are inherited from a second-pass Undermind synthesis (search conducted 2026-05-21); the primary source has *not* been verified at landing time. A primary-source verification spike is queued — see Working Notes above.

<i>ASF concern</i>	<i>Prior-art language</i>	<i>Relationship / Positioning</i>
Network-reconstruction limits from temporal data	Angulo-Moreno-Barabási-Liu 2015, <i>Fundamental limitations of network reconstruction</i> (arXiv:1508.03559) ⁶	<i>adjacent literature</i> — graph topology is generically as hard to recover as the full interaction matrix without persistent excitation; underdetermines coupling from per-node observation. Closest network-level impossibility neighbor on the topology-from-component-data axis. Different setting (general dynamical, not strategic), but the underdetermination-from-reduced-observation pattern is shared
Delay-reconstruction recovers only transitive closure	Cummins-Gedeon-Spendlove 2015, <i>On the Efficacy of State Space Reconstruction Methods in Determining Causality</i> (SIAM J. Appl. Dyn. Syst. 14:335) ⁷	<i>adjacent literature</i> — componentwise trajectory data recovers at most the transitive closure of the interaction graph, not direct edges or feedback orientation. Explicit negative result on what trajectory data cannot tell an observer; not a game-structure witness but operationally similar shape
Game-side observational equivalence (quantal-response rationalization; consistent-game-set characterization)	Haile-Hortaçsu-Kosenok 2008, <i>On the Empirical Content of Quantal Response Equilibrium</i> (DOI 10.1257/AER.98.1.180) ⁸ (any interior action distribution rationalizable as QRE); Ziani-Chandrasekaran-Ligett 2016, <i>Efficiently Characterizing Games Consistent with Perturbed Equilibrium Observations</i> (DOI 10.7907/Z91Z42CF) ⁹ (characterizing the whole consistent set of games rather than point-identifying)	<i>adjacent literature</i> — stronger observational-equivalence statements on the game side than mere nonidentification, but at the equilibrium-rationalization layer rather than the per-component-marginal-trajectory layer this segment treats. Game-topology distinction (potential vs non-potential; cooperative vs cyclic) not the target of these results
Information-structure observational compatibility	Penalva-Ryall 2003, <i>Causal assessment in finite extensive-form games</i> ¹⁰ ; Penalva-Zuasti-Fabra-Ryall 2004, <i>Empirical Implications of Information Structure in Finite-Length Extensive-Form Games</i> ¹¹ ; Lehrer-Rosenberg-Shmaya 2006, <i>Signaling and mediation in Bayesian games</i> ¹²	<i>adjacent literature</i> — different information structures observationally compatible with respect to outcome / equilibrium distributions; another form of hidden structure surviving unchanged at the observable level. Different concern (information structure, signaling, mediation) but same shape (hidden distinction not visible at observable layer)

Novelty assessment. The cross-topology game witness with hidden distinction in potential-vs-non-potential or cooperative-vs-cyclic structure under matched per-component marginals appears uncoded relative to the prior-art search¹³; Witness 1 (R0-vs-R1 Cournot under matched marginal mean and variance) and Witness 2 (R0-vs-R2 matching-pennies under aggregation) are the specific instantiations the verdict identifies as uncoded. The closest structural analogues

⁶The paper and the characterization of its claims are inherited from a second-pass Undermind synthesis (search conducted 2026-05-21); the primary source has *not* been verified at landing time. A primary-source verification spike is queued — see Working Notes above.

⁷The paper and the characterization of its claims are inherited from a second-pass Undermind synthesis (search conducted 2026-05-21); the primary source has *not* been verified at landing time. A primary-source verification spike is queued — see Working Notes above.

⁸The paper and the characterization of its claims are inherited from a second-pass Undermind synthesis (search conducted 2026-05-21); the primary source has *not* been verified at landing time. A primary-source verification spike is queued — see Working Notes above.

⁹The paper and the characterization of its claims are inherited from a second-pass Undermind synthesis (search conducted 2026-05-21); the primary source has *not* been verified at landing time. A primary-source verification spike is queued — see Working Notes above.

¹⁰The paper and the characterization of its claims are inherited from a second-pass Undermind synthesis (search conducted 2026-05-21); the primary source has *not* been verified at landing time. A primary-source verification spike is queued — see Working Notes above.

¹¹The paper and the characterization of its claims are inherited from a second-pass Undermind synthesis (search conducted 2026-05-21); the primary source has *not* been verified at landing time. A primary-source verification spike is queued — see Working Notes above.

¹²The paper and the characterization of its claims are inherited from a second-pass Undermind synthesis (search conducted 2026-05-21); the primary source has *not* been verified at landing time. A primary-source verification spike is queued — see Working Notes above.

¹³The paper and the characterization of its claims are inherited from a second-pass Undermind synthesis (search conducted 2026-05-21); the primary source has *not* been verified at landing time. A primary-source verification spike is queued — see Working Notes above.

are Bielecki et al. 2011 (multivariate Markov-chain same-marginals-different-coupling) and the Courbage-Hamdan / Hamdan sequence (marginal-non-determining stochastic processes) — both at the stochastic-process layer rather than the game-topology layer. AAT’s contribution: the *naming claim* on the game-topology-bit witness as the cross-regime indistinguishability instantiation of the established marginal-non-determining shape.

Search Log:

- 2026-05-21 (*targeted*): The core citations in §”Witness 1” / §”Witness 2” (Liberzon 2003, Dayawansa-Martin 1999, Shorten et al. 2007, Monderer-Shapley 1996, Hart-Mas-Colell 2000) are inherited from existing AAT machinery’s citation set; the witness construction is a parameter-matched application of standard game-theoretic templates against R0 shared-objective dynamics — the constructions are not novel in themselves, and the contribution is the cross-regime marginal-indistinguishability *recognition* and the mechanism-reduction to Instance 3’s rank-collapse framework.
- 2026-05-21 (*targeted second-pass*): A second-pass comprehensive-search pass searched specifically for prior art on (Q1) broader-topology-bit unidentifiability beyond the coupling-sign-bit witness, and (Q2) convergence-rate-class as a passive regime classifier. Verdict on both: *no exact hit found in this set*. The Related Work table above incorporates the verdict’s strongest neighbors at adjacent layers; see footnote [^bg2-2026-05-21] for the verification-deferral notice on every neighbor citation introduced from this pass.

Working Notes.

- **Provenance.** Derivation worked in spikes/.integrated/spike-identifiability-floor-instance-6-2026-05-21.md §3.3. The spike’s §4 M1 distinctness check identified three of four tests failing against Instance 3, landing the recommendation as “broaden Instance 3 + new appendix segment” rather than “new Instance 6.” This segment is the new appendix segment per that landing.
- **Cross-regime witness sub-cases.** Witness 1 (R0 vs R1) is the load-bearing cross-regime construction. Witness 2 (R0 vs R2) is qualitatively cleaner (the joint regime differs more dramatically) but operationally similar. A fourth witness (R1 vs R2 — Cournot equilibrium with parameter-matched marginals to matching-pennies cyclic dynamics) is plausible but not derived here; per the strengthen-first second-pass, the R0 vs R1 and R0 vs R2 witnesses are sufficient to establish the broadened no-go scope.
- **Connection to escape (e).** The convergence-rate-class escape (per #disc-identifiability-floor Instance 3 escape (e), added 2026-05-21) distinguishes Witness 2’s R0-vs-R2 case passively: Composite A converges exponentially, Composite C converges polynomially. Escape (e) does *not* distinguish Witness 1’s R0-vs-R1 case (both converge exponentially at the matched local rate α); the R0-vs-R1 distinction requires one of escapes (a)-(d) per the Instance-3 Boundary Characterization.
- **Composite-of-composites lift.** Per #disc-dynamic-regime-axis §I (Q3), the regime axis under nested composition is open. A natural extension of this segment’s witnesses to composite-of-composites would test whether marginal-indistinguishability at the sub-composite layer propagates to the meta-composite regime question. Follow-on spike candidate if the composition-of-regimes question becomes load-bearing for 04-eli-core/ population work.
- **Fano-style finite-sample refinement.** The §”Witness 1” / §”Witness 2” constructions are exact-population statements. A finite-sample bound on regime-classification minimax error at observation budget T would be the natural quantitative companion — per the 2026-05-20 spike’s Fano-anchor-as-finite-sample-refinement distinction. Plausibly derivable from #der-agent-opacity H_b machinery applied at the composite layer; not in scope here.
- **Migration note (2026-05-09 GUC rename):** Class 2 $\leftrightarrow$ Class 3 swap. Pre-2026-05-09: Class 2 = fully merged, Class 3 = partially modular. Post: Class 2 = Partial, Class 3 = Coupled. This segment uses regime-axis vocabulary (R0/R1/R2) rather than GUC-class vocabulary; no GUC-class references inside. Removed at candidate stage per FORMAT.md Gate 4.
- **Primary-source verification spike queued (2026-05-21).** All BG2-introduced citations in this segment’s Related Work table (marked with footnote [^bg2-2026-05-21]) need primary-source verification before any further reliance. The characterizations attributed to those papers are currently synthesized from ref/Unidentifiability_and_rate_class_prior_art.md,

not from direct reading of the sources. Verification targets in approximate load-bearing order for this segment: (i) Bielecki-Jakubowski-Niewkeglowski 2011 — *closest structural analogue* (multivariate-Markov-chain same-marginals-different-coupling); load-bearing for the segment's positioning of the Witness math; (ii) Courbage-Hamdan 2001 + Courbage-Hamdan 1998 + Hamdan 2000 + Hamdan 2002 sequence — strongest direct-precedent for “finite-order marginals do not determine the generating law”; the Hamdan 2000 journal name was not specified in the Undermind report and should be resolved at verification (the DOI prefix would identify it, but the executor will not infer from DOI structure without confirming the actual paper); (iii) Angulo et al. 2015, Cummins et al. 2015 — network-reconstruction limits; (iv) Haile et al. 2008, Ziani et al. 2016 — game-side observational equivalence; (v) Penalva et al. 2003/2004, Lehrer et al. 2006 — information-structure observational compatibility. Spike to be scheduled per Joseph's go.

Appendix AO

Derivation: Strategy-Proofness Impossibility (Gibbard-Satterthwaite)

The first charter instance of #disc-implementation-impossibility: Gibbard 1973 and Satterthwaite 1975 forbid an external designer from constructing a dominant-strategy incentive-compatible (DSIC), non-dictatorial, surjective social-choice function over ≥ 3 alternatives under unrestricted preferences. The AAT-side translation locates the designer's task in the composite-agent setting (#scope-composite-agent), recognizes the agents as composite sub-agents reporting preferences derived from their objectives $O_i^{(i)}$, and maps the four standard escapes against AAT machinery — with the preference-domain-restriction escape (single-peaked / single-crossing) named precisely as *adjacent to but not identical with* the sub-scope α' machinery of #deriv-strategic-composition. The adjacency is not identity: sub-scope α' secures best-response *convergence* under fixed reports; GS forbids *strategy-proof revelation*. The honest map of where AAT meets mechanism design is the contribution; the imported theorem is classical.

Formal Expression.

Setting. Let A be a finite alternative set with $|A| \geq 3$, and let L denote the set of strict linear orderings on A . A **social-choice function** is a map $f : L^n \rightarrow A$ that aggregates n agents' reported preferences $R = (R_1, \dots, R_n) \in L^n$ into a single outcome $f(R) \in A$.

Three constraints define the designer's regime:

Definition (dsic, strategy-proofness / dominant-strategy incentive compatibility).

$$\forall i, \forall R_{-i} \in L^{n-1}, \forall R_i, R'_i \in L : f(R_i, R_{-i}) R_i f(R'_i, R_{-i}). \quad (\text{AO.1})$$

Truth-telling is a dominant strategy: no agent i can profit from misreporting given any reports by the others.

Definition (non-dictatorial).

$$\neg \exists i^* : \forall R \in L^n, f(R) \text{ is } i^* \text{'s top-ranked alternative in } R_{i^*}. \quad (\text{AO.2})$$

No agent fully determines the outcome.

[*Definition (surjective)*] $f(L^n) = A$. Every alternative is reachable from some report profile.

External theorem.

Theorem (Gibbard 1973; Satterthwaite 1975). If $|A| \geq 3$ and preferences are unrestricted (each R_i ranges over all of L), then no social-choice function $f : L^n \rightarrow A$ is simultaneously DSIC, non-dictatorial, and surjective.

Postulate (gibbard-satterthwaite-1973-75, imported), exact.

References. Gibbard, A. (1973), “Manipulation of voting schemes: a general result,” *Econometrica* 41(4):587–601. Satterthwaite, M. A. (1975), “Strategy-proofness and Arrow’s conditions: existence and correspondence theorems for voting procedures and social welfare functions,” *J. Econ. Theory* 10(2):187–217. Modern textbook synthesis: Mas-Colell, Whinston & Green (1995), *Microeconomic Theory*, §23.C; geometric/topological proof structure shared with Arrow’s theorem in Reny, P. J. (2001), “Arrow’s theorem and the Gibbard-Satterthwaite theorem: a unified approach,” *Econ. Letters* 70:99–105.

No-go. A designer who insists simultaneously on DSIC, non-dictatorial, surjective, and unrestricted-preferences cannot construct the social-choice function. Any candidate f satisfying three of the four conditions violates the fourth.

AAT-side translation.

Within AAT’s composite-agent setting, the GS no-go translates as follows:

Derived (aat-mechanism-translation, from ‘#scope-composite-agent’ + ‘#deriv-strategic-composition’), conditional on the composite-revelation-mechanism reading.

- The “agents” are *sub-agents* in a composite per #scope-composite-agent, each carrying its own objective $O_t^{(i)}$ inducing a preference ordering $R_i \in L$ over composite-level alternatives A (the joint outcomes reachable through the composite’s strategic-composition machinery).
- A “social-choice function” $f : L^n \rightarrow A$ is a *mechanism* the designer (external to the composite agent) imposes — a rule that takes the sub-agents’ reports $R = (R_1, \dots, R_n)$ and produces a composite-level outcome. The designer’s task is to choose f .
- “Strategy-proofness” requires that each sub-agent’s truth-telling about its own $O_t^{(i)}$ -induced ordering is its best-response strategy. The DSIC condition is stronger than Nash equilibrium: dominance against *all* other-agent reports, not just equilibrium-consistent ones.

The GS theorem then states: under the four stated conditions, no such mechanism exists. The designer faces a structural impossibility on what can be implemented over the composite of strategically-reporting sub-agents under unrestricted preference domain.

Boundary characterization.

Table AO.1

<i>Escape</i>	<i>Mechanism-design route</i>	<i>AAT machinery?</i>
Preference-domain restriction	Single-peaked preferences (Black 1948 <i>J. Polit. Econ.</i> 56:23–34) admit the median-voter rule; single-crossing extends this. The domain restriction circumvents the unrestricted-preferences condition.	Adjacent to AAT machinery. Sub-scope α' of #deriv-strategic-composition (potential / monotone game structure) is <i>adjacent</i> to single-peaked / single-crossing — both restrict the preference-or-utility structure to admit convergence/aggregation guarantees — but not identical (see <i>Sub-scope α' vs preference-domain restriction</i> below).
Weakened solution concept	Bayes-Nash incentive compatibility (BIC) instead of dominant-strategy IC — d'Aspremont & Gérard-Varet (1979) “Incentives and incomplete information,” <i>J. Public Econ.</i> 11:25–45. The designer relaxes from dominance to expected-utility best-response under a common prior on others’ types.	Outside AAT machinery. BIC requires a shared prior over types and Bayesian reasoning over expected reports; AAT’s strategic-composition machinery does not currently formalize the mechanism-induced report game in BIC terms.
Randomization	Gibbard (1977) “Manipulation of schemes that mix voting with chance,” <i>Econometrica</i> 45:665–681 — random-dictator and convex-combination mechanisms admit strategy-proofness in expectation.	Outside AAT machinery. AAT’s machinery treats deterministic strategies Σ_i ; randomized-mechanism design is not in canon.
Restricted strategy space (enforced truth-telling)	The revelation-principle-companion: if the designer can verify reports against ground truth (e.g., observable types), strategy-proofness becomes definitional rather than incentivized.	Outside AAT machinery. AAT’s agent-spectrum (#def-agent-spectrum) does not currently include verification-of-self-reports machinery.

Strengthened consequence. The boundary-precise contribution: the AAT/mechanism-design intersection is at *one* of the four GS escapes, and that intersection is *adjacent without being identical*. The honest map looks like this.

Sub-scope α' vs preference-domain restriction. Three reframings under which sub-scope α' would *be* the GS escape all fail: sub-scope α' secures best-response *convergence* of the composite-dynamics given a fixed mechanism, not *strategy-proofness* of the mechanism itself. Concretely: agents in a potential game (the canonical sub-scope α' instance) can have profitable unilateral deviations under non-truthful preference revelation even when the joint best-response dynamics under the current revelation converge. The convergence-under-current-reports guarantee that sub-scope α' supplies is *not* the dominance-over-all-reports guarantee DSIC requires.

The honest read: sub-scope α' is *adjacent* to the preference-domain-restriction escape because both restrict the preference-or-payoff structure in similar ways (single-peaked preferences exhibit a potential-game-like ordinal structure on single-dimensional issue spaces; monotone games on lattices share order-theoretic features with single-crossing preference domains), but it is *not the same restriction*. Sub-scope α' ’s contribution to the composite-agent layer is convergence-given-mechanism, not mechanism-strategy-proofness. The boundary is named precisely; AAT does not claim to escape GS, but recognizes where the framework’s strategic-composition machinery sits *next* to one of GS’s classical escapes.

The other three escapes belong entirely to mechanism design and social-choice theory. The framework honestly documents that, and locates its current contribution as *recognition of the boundary* rather than *provision of the escapes*.

Epistemic Status.

Exact for the imported theorem (Gibbard 1973, Satterthwaite 1975 — classical theorems with closed-form proofs in the source literature, primary-source verification scheduled for promotion past draft).

Conditional for the AAT-side translation — the mapping from social-choice-function to composite-mechanism-over-sub-agent-reports is mechanical under the *composite-revelation-mechanism reading* (sub-agents per #scope-composite-agent; reports as inputs to a designer-chosen f ; composite outcomes as elements of A), but the reading itself is a representational choice; alternative AAT readings (e.g., treating sub-agents as auctioning under VCG-like settings; treating f as a Σ -level commitment device) exist and would map GS to different AAT-side vocabulary.

Robust qualitative for the boundary characterization (sub-scope α' is adjacent-but-not-identical to preference-domain restriction; the other three escapes are outside AAT canon) — survives across instance details, with the specific functional adjacency (single-peaked $\leftrightarrow$ potential-game-like structure on issue spaces) noted but not derived as identity.

Max attainable: *Exact* for the imported theorem; *derived conditional* for the AAT translation; *robust qualitative* for the boundary characterization. The AAT translation's max ceiling stays *conditional* unless an alternative reading is forced by some AAT-internal uniqueness argument (no such argument currently in hand).

Discussion.

Why this is not a #disc-identifiability-floor instance. GS fails the identifiability-floor's five-element shape on actor positioning: in Instances 1–4 of the floor the actor frustrated by the no-go is the *agent itself* (cannot detect, identify, certify, or recover an architectural d.o.f. from limited data), and the escape elevates *agent-side* information-augmentation machinery (loop-interventional access, latent observability, matched-Tier composite-extended observation, similarity-fiber resolution). GS's actor is the *designer*; the escapes are *design-constraint relaxations*. The cluster is structurally different and belongs in the sister meta-segment #disc-implementation-impossibility. This segment is GS's home there.

Sub-scope α' adjacency, named precisely. The Discussion above states sub-scope α' is *adjacent* to preference-domain restriction but *not identical*; this is load-bearing for the honest scope-marking discipline of #disc-implementation-impossibility. The framework's contribution here is recognition of where the strategic-composition machinery sits next to one of GS's classical escapes — not a claim that AAT supplies the escape. Mechanism-design proper supplies the escape; AAT documents the proximity.

Peer voice with the mechanism-design literature. The imported theorem is classical; the proof structure (Gibbard 1973 via the revelation-principle-style argument; Satterthwaite 1975 via direct manipulation construction; Reny 2001 via geometric/topological unification with Arrow) is mature. AAT does not re-prove the theorem; AAT does not displace the standard escape routes; AAT does not claim that the strategic-composition machinery solves mechanism design. The segment's contribution is honest recognition of the boundary, with the imported theorem cited under its original names and full bibliographic references.

Relationship to #deriv-strategic-composition. The placement of “mechanism-design impossibility” as a candidate-fourth-instance of #disc-identifiability-floor is wrong on actor-positioning grounds. GS's home is here in the sister meta-segment; sub-scope α' retains its load-bearing role in strategic-composition for best-response convergence, and is *additionally* recognized as adjacent to GS's preference-domain-restriction escape under the boundary characterization above.

Findings.

The Gibbard-Satterthwaite Translation and Its Sub-Scope α' Adjacency. **Brief:** The Gibbard-Satterthwaite theorem says that no voting rule can be at once strategy-proof (no one ever benefits from misreporting their preferences), non-dictatorial (no single voter always decides), surjective (every alternative is reachable), and broadly applicable (works on unrestricted preferences) when there are at least three alternatives. The classical escapes have been studied for decades — restrict the preference domain (require single-peaked preferences; the median voter rule works), allow randomization (random dictator is strategy-proof in expectation), accept a weaker incentive concept (Bayes-Nash IC under a shared prior), or constrain the strategy space (enforce truth-telling by verifying reports). The contribution here is the honest map: when the framework recognizes a composite agent whose sub-agents have their own objectives, the strategic-composition machinery’s sub-scope α' (potential or monotone games — the regime where best-response dynamics converge) is *adjacent to* the preference-domain-restriction escape (single-peaked / single-crossing) — both restrict the preference structure to admit aggregation/convergence guarantees — but is *not the same restriction*: convergence-under-current-reports is not dominance-over-all-reports. The other three classical escapes belong entirely to mechanism design and social-choice theory; the framework documents them as outside the current canon. Mapping that boundary honestly is the contribution; the impossibility is classical.

Impact: Closes the candidate-fourth-instance flag in #deriv-strategic-composition Discussion by landing GS in its proper home — the designer-side #disc-implementation-impossibility meta-segment, with sub-scope α' ’s relationship to the preference-domain-restriction escape named precisely. The named adjacency-without-identity is the kind of scope-honesty discipline the framework runs throughout: AAT recognizes where its machinery sits next to classical mechanism-design escape routes, and refuses the temptation to overclaim that sub-scope α' “supplies” the escape (it does not — it supplies a different guarantee at a different layer of the composite-agent dynamics). The result is a charter instance of the designer-side cluster with the AAT/mechanism-design boundary mapped with precision rather than enthusiasm.

Novelty Claim: *Claim recognition* of the AAT-side translation of the Gibbard-Satterthwaite theorem into the composite-agent setting (sub-agents per #scope-composite-agent; reports as inputs to a designer-chosen mechanism; composite outcomes as the alternative set); *claim differentiation* on the sub-scope α' $\leftrightarrow$ preference-domain-restriction adjacency: sub-scope α' is *adjacent to but not identical with* the GS preference-domain escape (the three-reframing strengthen-first check is the documented argument). The imported theorem is classical (Gibbard 1973; Satterthwaite 1975); the AAT-side translation, the boundary-characterization annotations, and the honest adjacency-without-identity recognition are the contributions.

Related Work:

Table AO.2

<i>ASF concern</i>	<i>Prior-art language</i>	<i>Relationship / Positioning</i>
The imported theorem	Gibbard, A. (1973), “Manipulation of voting schemes: a general result,” <i>Econometrica</i> 41(4):587–601 (found pre-2026); Satterthwaite, M. A. (1975), “Strategy-proofness and Arrow’s conditions: existence and correspondence theorems for voting procedures and social welfare functions,” <i>J. Econ. Theory</i> 10(2):187–217 (found pre-2026)	<i>formal antecedent</i> — primary-source verification scheduled for promotion past draft
Modern textbook synthesis	Mas-Colell, A., Whinston, M. D. & Green, J. R. (1995), <i>Microeconomic Theory</i> , Oxford University Press, §23.C (found pre-2026)	<i>adjacent literature</i> — used for the constraint-regime formulation in the Setting above

ASF concern	Prior-art language	Relationship / Positioning
Geometric/topological proof unification with Arrow	Reny, P. J. (2001), “Arrow’s theorem and the Gibbard-Satterthwaite theorem: a unified approach,” <i>Econ. Letters</i> 70:99–105 (found 2026-05-20)	<i>adjacent literature</i> — names the topological-combinatorial mechanism shared with Arrow’s theorem (Instance 3 of #disc-implementation-impossibility); informs the mechanism-family discussion in the meta-segment
Single-peaked preferences as the canonical domain-restriction escape	Black, D. (1948), “On the rationale of group decision-making,” <i>J. Polit. Econ.</i> 56:23–34 (found pre-2026)	<i>adjacent literature</i> — the canonical escape; the sub-scope α' adjacency is named relative to this
Random-dictator and randomized-mechanism escapes	Gibbard, A. (1977), “Manipulation of schemes that mix voting with chance,” <i>Econometrica</i> 45:665–681 (found pre-2026)	<i>adjacent literature</i> — outside AAT machinery; documented as honest scope-mark
Bayes-Nash IC weakening	d’Aspremont, C. & Gérard-Varet, L.-A. (1979), “Incentives and incomplete information,” <i>J. Public Econ.</i> 11:25–45 (found pre-2026)	<i>adjacent literature</i> — outside AAT machinery; documented as honest scope-mark

Search Log:

- 2026-05-22 (*intuition-only* on the AAT-side composite-revelation-mechanism reading): The mapping from social-choice function to composite-mechanism-over-sub-agent-reports (sub-agents per #scope-composite-agent; reports as inputs; composite outcomes as A) is the natural AAT-side reading; alternative readings exist (treating f as a Σ -level commitment device; treating sub-agents as VCG-auction bidders) and would map GS to different AAT-side vocabulary. The composite-revelation reading is chosen here for parsimony with #deriv-strategic-composition; targeted future search candidates would be alternative-reading derivations if a downstream use of this segment requires them.

Working Notes.

- **Provenance.** Authored 2026-05-22 as commit 2 of the post-Track-CR Track-B cycle, executing spikes/implementation-impossibility-meta-segment-plan.md §4.a. The 2026-05-20 strengthen-first arm record at spikes/.integrated/spike-4th-identifiability-floor-instance-2026-05-20.md §4 carries the load-bearing argument that sub-scope α' is *adjacent to but not identical with* the GS preference-domain-restriction escape: it re-derived GS against the identifiability-floor’s five-element test, tested three reframings under strengthen-before-soften (GS as identifiability-of-truthful-preferences; GS as Part III multi-agent no-go on #scope-composite-agent (C-iv) routes; GS as alignment-design obstruction), and confirmed sub-scope α' does *not* escape GS (it secures best-response convergence, not strategy-proofness). The verdict’s recommended disposition — route GS to a sister meta-segment if mechanism design becomes a first-class framework concern — is what this segment lands.
- **Open before candidate.** (a) Primary-source verification of Gibbard 1973 and Satterthwaite 1975 — the theorem statement above is taken from the textbook synthesis (Mas-Colell, Whinston & Green 1995 §23.C) and the geometric-proof-unification paper (Reny 2001); direct primary-source reading is queued. (b) Alternative AAT-side readings of GS — treating f as a Σ -level commitment device, or treating sub-agents as VCG-auction bidders — would map GS into different AAT vocabulary; whether either alternative reading produces a stronger AAT-side derivation is a candidate sub-spike. (c) The sub-scope $\alpha' \leftrightarrow$ single-peaked-preferences adjacency is currently *robust qualitative* — there may be an AAT-internal derivation showing the structural correspondence at parameter-level (e.g., that single-peaked preferences on an issue-space induce a potential-game structure under standard best-response dynamics); the derivation is queued and would lift the adjacency claim to *exact* if it lands cleanly. Indexed at spikes/PROPOSED.md Tier 3 (added 2026-05-24) as a segment-perspective audit-judgment with this segment as canonical home; the lift is symmetric for #deriv-social-welfare-aggregation-impossibility (Arrow’s preference-domain-restriction escape shares the same adjacency).

- **The composite-revelation-mechanism reading and its sub-agents.** The AAT-side translation treats the GS “agents” as sub-agents in a composite per `#scope-composite-agent`, each carrying its own objective $O_t^{(i)}$ inducing a preference ordering over composite-level alternatives. This is the most parsimonious AAT-side reading and aligns with `#deriv-strategic-composition`’s framing of strategic composition over sub-agent objectives. The reading does *not* assume sub-agents have unbounded computational capacity, do not constrain their reporting strategies to be honest, do not require shared priors — the GS no-go applies to the designer’s mechanism task regardless of those details.
 - **The cluster’s discipline filter.** GS is the *cleanest* charter instance of `#disc-implementation-impossibility` — the discipline filter (constructive-impossibility-shape only) admits it directly: the setting is well-stated, the external theorem is classical, the no-go is sharp, the escapes are well-characterized in the literature, and the strengthened consequence (the AAT/mechanism-design boundary precisely mapped at the sub-scope α' adjacency) is the framework’s honest contribution. The next charter instances (MS, Arrow) fit the same template with their own variations.
-

Appendix AP

Derivation: Bilateral Trade Impossibility (Myerson-Satterthwaite)

The second charter instance of #disc-implementation-impossibility: Myerson & Satterthwaite 1983 forbid an external designer from constructing an ex-post efficient, Bayesian incentive compatible, individually rational, budget-balanced bilateral-trade mechanism when buyer and seller have private valuations drawn from overlapping continuous distributions. The AAT-side translation locates the buyer and seller as a two-sub-agent composite per #scope-composite-agent; recognizes that the two-agent strategic-composition machinery of #deriv-strategic-composition applies cleanly to the agents' best-response dynamics *given* a fixed mechanism but does not cover the designer's mechanism-choice problem; and honestly maps the four classical escapes as almost entirely outside AAT's current canon. The contribution here is *boundary precision*: bilateral-trade impossibility is *adjacent to* AAT (the strategic-composition layer engages the agents-given-mechanism dynamics) but the designer's task lives in mechanism-design proper. The honest scope-mark is the contribution.

Formal Expression.

Setting. Two parties: a buyer B with private valuation $v_b \in [0, \bar{v}_b]$ drawn from distribution F_b with density f_b , and a seller S with private valuation $v_s \in [0, \bar{v}_s]$ drawn from distribution F_s with density f_s . The supports overlap: there exists a non-empty interval $[a, b] \subseteq [0, \bar{v}_b] \cap [0, \bar{v}_s]$ such that f_b and f_s are both strictly positive on $[a, b]$. A **direct mechanism** is a pair (p, x) where $p : [0, \bar{v}_b] \times [0, \bar{v}_s] \rightarrow [0, 1]$ is the trade probability and $x : [0, \bar{v}_b] \times [0, \bar{v}_s] \rightarrow \mathbb{R}$ is the buyer's expected payment to the seller (with the seller receiving x). Equivalently with interim allocation rules $\bar{p}_b(v_b)$, $\bar{p}_s(v_s)$ and interim transfer rules $\bar{x}_b(v_b)$, $\bar{x}_s(v_s)$ for buyer-and-seller respectively.

Four constraints define the designer's regime:

[Definition (*ex-post-efficient*)] $p(v_b, v_s) = 1 \iff v_b \geq v_s$. Trade occurs exactly when the buyer values the good at least as much as the seller.

[Definition (*bayesian-ic, BIC*)] Each agent's truthful reporting is a Bayesian best response given the other's truthful strategy:

$$v_b \bar{p}_b(v_b) - \bar{x}_b(v_b) \geq v_b \bar{p}_b(\hat{v}_b) - \bar{x}_b(\hat{v}_b) \quad \forall v_b, \hat{v}_b, \quad (\text{AP.1})$$

and symmetrically for the seller (with the sign reversed since the seller is paid).

[Definition (*individual-rationality, IR*)] Both parties participate voluntarily: expected interim payoff is nonnegative for every type:

$$v_b \bar{p}_b(v_b) - \bar{x}_b(v_b) \geq 0 \quad \forall v_b, \quad \text{and} \quad \bar{x}_s(v_s) - v_s \bar{p}_s(v_s) \geq 0 \quad \forall v_s. \quad (\text{AP.2})$$

[Definition (budget-balance)] No external subsidy: the buyer's expected payment exactly equals the seller's expected receipt:

$$\mathbb{E}[x(v_b, v_s)] = 0 \quad (\text{net designer-side}). \quad (\text{AP.3})$$

External theorem.

Theorem (Myerson & Satterthwaite 1983). If f_b and f_s are strictly positive on a common open interval and the supports overlap, no direct mechanism (p, x) is simultaneously ex-post efficient, BIC, IR, and budget-balanced.

Postulate (myerson-satterthwaite-1983, imported), exact.

Reference. Myerson, R. B. & Satterthwaite, M. A. (1983), "Efficient mechanisms for bilateral trading," *J. Econ. Theory* 29(2):265–281. Modern textbook synthesis: Krishna, V. (2002), *Auction Theory*, Academic Press, §5.3; Mas-Colell, A., Whinston, M. D. & Green, J. R. (1995), *Microeconomic Theory*, Oxford University Press, §23.E.

No-go. A designer demanding ex-post efficiency + BIC + IR + budget-balance cannot construct the bilateral-trade mechanism. Any candidate mechanism satisfying three of the four conditions violates the fourth. Equivalently: under private values with overlapping supports, the designer must accept a deficit (relax budget-balance), accept inefficiency (relax ex-post efficiency), accept that participation is not always individually rational, or weaken the incentive concept.

AAT-side translation.

Within AAT's composite-agent setting, the MS no-go translates partially:

Derived (aat-bilateral-trade-translation, from '#scope-composite-agent'), conditional on the buyer-seller-as-two-sub-agent reading, partial.

- The "buyer" and "seller" are *sub-agents* in a two-sub-agent composite per #scope-composite-agent, each carrying its own objective $O_t^{(i)}$ — the buyer's $O_t^{(B)}$ rewarding acquisition at favorable price; the seller's $O_t^{(S)}$ rewarding relinquishing at favorable price.
- The "mechanism" (p, x) is a designer-imposed structure on the composite's joint action space and reward redistribution — outside AAT's agent-level machinery; AAT's strategic-composition machinery in #deriv-strategic-composition applies to the agents' best-response dynamics *given* the mechanism (p, x) , not to the designer's choice of (p, x) .

The honest translation: the *agents-given-mechanism layer* is internal to AAT (strategic-composition machinery applies); the *mechanism-choice layer* is external. The MS no-go applies at the mechanism-choice layer; AAT does not currently cover it.

Boundary characterization.

Table AP.1

<i>Escape</i>	<i>Mechanism-design route</i>	<i>AAT machinery?</i>
Subsidies (relax budget-balance)	VCG mechanism with deficit; pivotal-mechanism transfers. The designer covers any shortfall from external resources.	Outside AAT machinery. AAT does not currently formalize designer-side transfer pools or subsidy budgets.
Second-best mechanisms (relax ex-post efficiency)	Myerson 1981 “Optimal auction design,” <i>Math. Oper. Res.</i> 6:58–73 — characterizes revenue-optimal mechanisms with inefficient withholding in the marginal-cost-curve region; bilateral-trade analog accepts trade-failure when $v_b - v_s$ is small relative to the inefficiency cost.	Outside AAT machinery. Optimal-mechanism design under welfare-relaxation is the constructive side of mechanism design; deliberately excluded by the discipline filter (see #disc-implementation-impossibility’s “constructive mechanism-design” exclusion).
Common values (relax private values)	Information aggregation under shared-value components — Wilson 1977 “A bidding model of perfect competition,” <i>Rev. Econ. Stud.</i> 44:511–518; the no-go softens because the agents’ reports carry shared signal.	Outside AAT machinery; adjacent to #scope-edge-update-causal-validity’s regime-indexed identification structure. The shared-signal regime is structurally similar to AAT’s partial-intervention Regime B (#scope-edge-update-causal-validity’s middle tier), but the formal connection has not been derived.
Long-run reputation / repeated interaction	Single-shot MS does not bind in indefinitely-repeated bilateral-trade games — Folk-theorem-style mechanisms with reputation backlash. Fudenberg-Maskin 1986 <i>Econometrica</i> 54:533–554 for the canonical Folk-theorem setup.	Outside AAT machinery; adjacent to the strategic-composition machinery in repeated-game form. The repeated-game extension of #deriv-strategic-composition is open canonical work; the MS-escape connection would land there if pursued.

Strengthened consequence. The MS boundary lies *almost entirely outside* AAT’s current machinery — and the *honest delineation* is the contribution. AAT does not claim to escape MS; AAT does not provide subsidies, optimal-mechanism design, common-value aggregation, or repeated-game reputation machinery in its current canon. The framework’s contribution at this charter instance is precisely the boundary recognition: the two-agent strategic-composition machinery in #deriv-strategic-composition applies cleanly to the buyer-seller best-response dynamics *given* a fixed mechanism (the agents-given-mechanism layer is internal), but the designer’s mechanism-choice problem (the mechanism-choice layer) is external. The two escapes flagged as *adjacent* (common-values to Regime B; repeated-interaction to repeated-strategic-composition) are open avenues if AAT extends in those directions; the two flagged as *outside* (subsidies, second-best mechanisms) belong to mechanism-design proper.

The contribution is the same kind of scope-honesty discipline #disc-identifiability-floor runs on the agent side and #deriv-strategy-proofness-impossibility runs at a different AAT-machinery-intersection: name the boundary precisely; refuse to overclaim machinery the framework does not currently supply.

Epistemic Status.

Exact for the imported theorem (Myerson & Satterthwaite 1983 — classical theorem with closed-form proof in the source literature; primary-source verification scheduled for promotion past draft).

Conditional, partial for the AAT-side translation — the buyer-seller-as-two-sub-agent reading is the natural composite-agent translation, but the *agents-given-mechanism layer* is the only part of the no-go internal to AAT canon; the *mechanism-choice layer* (where the no-go binds) is external. The partial-translation is honest — AAT recognizes the bilateral-trade designer-task as adjacent but does not formalize the designer-side mechanism-choice problem.

Robust qualitative for the boundary characterization — four escapes documented with honest scope-marks; two adjacencies named (common-values $\leftrightarrow$ Regime B; repeated-interaction $\leftrightarrow$ repeated-strategic-composition); no precise functional intersection derived.

Max attainable: *Exact* for the imported theorem. *Derived conditional* for the AAT translation, with the *conditional* gating on the buyer-seller-two-sub-agent reading; an alternative reading (e.g., buyer and seller as *single-agent-with-two-roles* under intertemporal preference modeling) would change the translation but is not currently in canon. *Robust qualitative* for the boundary characterization; lifting to *exact* on the common-values $\leftrightarrow$ Regime B intersection would require formalizing the regime-correspondence, which is open.

Discussion.

Why this is not a #disc-identifiability-floor instance. Same actor-positioning argument as for the GS instance (#deriv-strategy-proofness-impossibility): the actor frustrated by the MS no-go is the *designer* (cannot construct a mechanism satisfying all four conditions), not an agent inferring from data. The escapes are *design-constraint relaxations* (relax budget-balance / efficiency / private-values / single-shot interaction), not information augmentations of an inferential regime. The cluster belongs in #disc-implementation-impossibility, not in the identifiability-floor.

Why the AAT translation is honestly partial. Unlike the GS instance (where sub-scope α' of #deriv-strategic-composition has a named adjacency to the preference-domain-restriction escape), MS has *no* AAT-internal escape adjacency. The two *adjacent* escapes (common-values $\leftrightarrow$ Regime B; repeated-interaction $\leftrightarrow$ repeated-strategic-composition) are open extensions, not current canon. The two *outside* escapes (subsidies, second-best mechanisms) belong to mechanism-design proper. The framework's contribution at this charter instance is therefore *purely the boundary recognition*: AAT covers the agents-given-mechanism layer, not the mechanism-choice layer; the MS no-go binds at the latter.

Peer voice with the mechanism-design literature. The imported theorem is classical; the proof structure (Myerson-Satterthwaite via Myerson's virtual-valuations machinery) is mature. AAT does not re-prove the theorem; AAT does not displace the standard escape routes; AAT does not claim that the strategic-composition machinery extends to mechanism-choice. The segment's contribution is honest recognition of the boundary, with the imported theorem cited under its original name and full bibliographic references.

Two open AAT-side extensions flagged. The common-values $\leftrightarrow$ Regime B adjacency and the repeated-interaction $\leftrightarrow$ repeated-strategic-composition adjacency are open avenues. If AAT extends in either direction — formalizing shared-signal information regimes alongside the current intervention-regime classification, or extending strategic-composition into repeated-game form — the MS instance's boundary characterization would gain AAT-machinery content where it currently has honest scope-marks. The flags here document the avenues without committing to either extension.

Findings.

The Bilateral Trade Boundary — Honest Scope-Marking as Contribution. **Brief:** When a buyer and seller each know their own valuation but not the other's, and the designer wants a trading mechanism that is at once fair (trade happens whenever the buyer values the good more than the seller does), incentive-honest (neither party can profitably misreport their valuation), voluntary (both parties want to participate), and self-funded (no outside subsidies), the Myerson-Satterthwaite theorem says all four cannot hold simultaneously. The classical escapes have been studied for decades — accept a budget deficit, accept inefficient withholding, exploit common-value information, or rely on long-run reputation in repeated trade. The honest contribution here is that *none* of these escapes is currently provided by the framework's machinery. The framework covers what the buyer and seller do *given* a mechanism (the two-sub-agent strategic-composition machinery applies cleanly to their best-response dynamics); the *designer's choice of mechanism* lives in mechanism-design proper, and the framework documents that boundary precisely rather than overclaiming reach. Two of the four classical escapes (common-values and repeated-interaction) are flagged as adjacent — open avenues if the framework extends in those directions — and two (subsidies and second-best mechanisms) belong entirely

to mechanism design. The contribution is the map of where the framework meets the bilateral-trade boundary, not the provision of an escape.

Impact: The cleanest example of the #disc-implementation-impossibility cluster’s discipline filter doing its honest work. Where the GS charter instance has a named sub-scope α' adjacency to make precise, the MS charter instance has *none* — the AAT-machinery contribution is purely boundary-recognition. That contrast within the cluster is itself epistemic content: the framework’s contribution is precision about *where its machinery contributes* (GS: adjacent to preference-domain restriction) and *where it does not* (MS: outside almost entirely). Two open avenues flagged (common-values $\leftrightarrow$ Regime B; repeated-interaction $\leftrightarrow$ repeated-strategic-composition) document where future framework extensions could intersect MS without committing to either.

Novelty Claim: *Claim recognition* of the AAT-side translation of the Myerson-Satterthwaite theorem under the buyer-seller-as-two-sub-agent reading; *claim differentiation* on the agents-given-mechanism layer being internal to AAT and the mechanism-choice layer being external — the no-go binds at the latter, which is the honest scope-mark. The imported theorem is classical (Myerson & Satterthwaite 1983); the AAT-side translation, the boundary-characterization annotations, the honest *almost-entirely-outside* scope-mark, and the two open adjacencies (common-values, repeated-interaction) are the contributions.

Related Work:

Table AP.2

<i>ASF concern</i>	<i>Prior-art language</i>	<i>Relationship / Positioning</i>
The imported theorem	Myerson, R. B. & Satterthwaite, M. A. (1983), “Efficient mechanisms for bilateral trading,” <i>J. Econ. Theory</i> 29(2):265–281 (found pre-2026)	<i>formal antecedent</i> — primary-source verification scheduled for promotion past draft
Modern textbook synthesis	Krishna, V. (2002), <i>Auction Theory</i> , Academic Press, \$5.3 (found pre-2026); Mas-Colell, Whinston & Green (1995), §23.E	<i>adjacent literature</i> — used for the constraint-regime formulation in the Setting above
Optimal mechanism design (second-best escape)	Myerson, R. B. (1981), “Optimal auction design,” <i>Math. Oper. Res.</i> 6(1):58–73 (found pre-2026)	<i>adjacent literature</i> — the canonical second-best-escape construction; outside AAT canon per the discipline filter exclusion
Common-values bidding (adjacent escape)	Wilson, R. (1977), “A bidding model of perfect competition,” <i>Rev. Econ. Stud.</i> 44(3):511–518 (found pre-2026)	<i>adjacent literature</i> — flagged as adjacent to #scope-edge-update-causal-validity Regime B; formal correspondence is open
Folk theorem (repeated-interaction escape)	Fudenberg, D. & Maskin, E. (1986), “The folk theorem in repeated games with discounting or with incomplete information,” <i>Econometrica</i> 54(3):533–554 (found pre-2026)	<i>adjacent literature</i> — flagged as adjacent to repeated-strategic-composition extension of #deriv-strategic-composition; formal correspondence is open

Search Log:

- 2026-05-22 (*intuition-only* on the buyer-seller-two-sub-agent reading): The natural composite-agent translation maps buyer and seller to two sub-agents per #scope-composite-agent with their own $O_t^{(i)}$. Alternative readings (single-agent-with-two-roles under intertemporal preference; large-economy aggregation limit) exist; not pursued here. Targeted future search candidates if alternative readings become relevant: Wilson 1977 large-economy bidding limit; behavioral / dual-self models of intertemporal trade.

- 2026-05-22 (*intuition-only* on the common-values $\leftrightarrow$ Regime B adjacency): The common-values escape's structural feature (agents' reports carry shared signal about a partially-observable joint state) is reminiscent of #scope-edge-update-causal-validity's Regime B (partial intervention; confounder-adjustment under side-channel observation). Whether the correspondence is formal or merely metaphorical is open. Targeted future search would compare the information-aggregation structure of common-values mechanisms (Wilson 1977; Pesendorfer-Swinkels 1997 *Econometrica* 65:1247–1281) against AAT's Regime B partial-intervention machinery.
- 2026-05-22 (*intuition-only* on the repeated-interaction $\leftrightarrow$ repeated-strategic-composition adjacency): The repeated-bilateral-trade escape's structure (reputation backlash sustaining cooperation through grim-trigger or carrot-and-stick strategies) is the natural extension of strategic-composition to repeated games. Whether the formal correspondence holds is open. Targeted future search would compare the Folk-theorem machinery (Fudenberg-Maskin 1986; Fudenberg-Levine-Maskin 1994 *Econometrica* 62:997–1039) against a repeated-game extension of #deriv-strategic-composition.

Working Notes.

- **Provenance.** Authored 2026-05-22 as commit 3 of the post-Track-CR Track-B cycle, executing spikes/implementation-impossibility-meta-segment-plan.md §4.b. The 2026-05-20 strengthen-first arm record at spikes/.integrated/spike-4th-identifiability-floor-instance-2026-05-20.md §5 names MS as a candidate instance held back from the floor cluster (along with Arrow) pending the Joseph-reserved structural-canon decision on a sister meta-segment — Joseph approved the sister-meta-segment landing in this cycle; MS is the second charter instance of #disc-implementation-impossibility.
 - **Open before candidate.** (a) Primary-source verification of Myerson & Satterthwaite 1983 — the theorem statement above is taken from the textbook synthesis (Krishna 2002 §5.3; Mas-Colell, Whinston & Green 1995 §23.E); direct primary-source reading is queued. (b) The two adjacent-escape correspondences (common-values $\leftrightarrow$ Regime B; repeated-interaction $\leftrightarrow$ repeated-strategic-composition) are *open* — flagged as adjacent in honest scope-mark, not derived. Lifting either to derived conditional would require formalizing the corresponding extension of AAT canon (#scope-edge-update-causal-validity Regime B's shared-signal interpretation, or a repeated-strategic-composition segment); both are open canonical work. (c) The buyer-seller-as-two-sub-agent reading is the natural composite-agent translation; alternative readings (single-agent-with-two-roles under intertemporal preference; large-economy aggregation) exist; not pursued here.
 - **The honest-almost-entirely-outside scope-mark.** MS is the cluster instance where the AAT/mechanism-design boundary is *most* outside AAT — even more so than GS (where sub-scope α' is named-adjacent) and likely more than Arrow (where cardinal preferences are adjacent to #form-objective-functional). The cluster's discipline filter admits MS *because* it fits the constructive-impossibility shape (well-stated setting, external classical theorem, sharp no-go, characterized escapes); the honest scope-mark on most escapes being outside AAT canon is the segment's contribution, not a weakness. The three charter instances together illustrate the cluster's range — GS with one named adjacency, Arrow with one named adjacency plus one open avenue, MS with no current adjacencies but two open avenues.
-

Appendix AQ

Derivation: Social Welfare Aggregation Impossibility (Arrow)

The third charter instance of #disc-implementation-impossibility: Arrow 1951 (expanded 1963) forbids an external designer from constructing a social welfare function aggregating n agents' strict preference orderings into a single social preference ordering over ≥ 3 alternatives under unrestricted domain (U), Pareto efficiency (P), Independence of Irrelevant Alternatives (IIA), and non-dictatorial (ND). The AAT-side translation locates the agents as composite sub-agents whose individual preference orderings are derived from their objectives $O_i^{(i)}$, and maps the four classical escapes with *two* AAT-machinery adjacencies named: the preference-domain-restriction escape (single-peaked / single-crossing $\rightarrow$ Condorcet-consistent rules) is *adjacent to* sub-scope α' of #deriv-strategic-composition (the same adjacency #deriv-strategy-proofness-impossibility names), and the cardinal-preference escape (utilitarianism, range voting) is *adjacent to* #form-objective-functional's value-functional machinery — both adjacencies open avenues, not elevations to load-bearing. The contribution: two AAT-side adjacencies precisely mapped, with the cardinal-preference adjacency unique to Arrow (it does not appear in GS or MS) because Arrow's ordinal-only constraint is what AAT's value-functional formulation could in principle relax while GS's strategy-proofness regime is orthogonal to the ordinal/cardinal distinction.

Formal Expression.

Setting. Let A be a finite alternative set with $|A| \geq 3$, and let L denote the set of strict linear orderings on A . A **social welfare function** is a map $F : L^n \rightarrow L$ that aggregates n agents' strict preference orderings $R = (R_1, \dots, R_n) \in L^n$ into a single social ordering $F(R) \in L$ on the same alternative set.

Four constraints define the designer's regime:

[*Definition (unrestricted-domain, U)*] The map F is defined on all of L^n — every possible profile of strict preference orderings is a valid input.

[*Definition (pareto-efficient, P)*] If every agent strictly prefers a over b (i.e., $a R_i b$ for all i), then the social ordering puts a strictly above b ($a F(R) b$).

[*Definition (independence-of-irrelevant-alternatives, IIA)*] The social ordering's comparison of any pair (a, b) depends only on the individuals' comparisons of (a, b) — not on their comparisons involving any third alternative c . Formally: for any two profiles R, R' such that each agent ranks a and b the same way in R and R' , the social orderings $F(R)$ and $F(R')$ rank a, b the same way.

[*Definition (non-dictatorial, ND)*] No agent i^* such that for every profile R and every pair a, b with $a R_{i^*} b$ strictly, the social ordering $F(R)$ also puts a above b — i.e., no agent's strict preferences are imposed on society regardless of others.

External theorem.

Theorem (Arrow 1951, 1963). If $|A| \geq 3$, then no social welfare function $F : L^n \rightarrow L$ is simultaneously U, P, IIA, and ND.

Postulate (arrow-1951-1963, imported), exact.

References. Arrow, K. J. (1951, 2nd ed. 1963), *Social Choice and Individual Values*, Yale University Press. The expanded 1963 edition adds the simplified proof and discusses extensions. Modern proof structures: Geanakoplos, J. (2005), “Three brief proofs of Arrow’s impossibility theorem,” *Econ. Theory* 26:211–215; Reny, P. J. (2001), “Arrow’s theorem and the Gibbard-Satterthwaite theorem: a unified approach,” *Econ. Letters* 70:99–105 (geometric/topological unification with GS).

No-go. A designer demanding U + P + IIA + ND cannot construct the social welfare function. Any candidate F satisfying three of the four conditions violates the fourth.

AAT-side translation.

Within AAT’s composite-agent setting, the Arrow no-go translates as follows:

Derived (aat-social-welfare-translation, from ‘#scope-composite-agent’ + ‘#form-objective-functional’), conditional on the composite-aggregation-mechanism reading.

- The “agents” are *sub-agents* in a composite per #scope-composite-agent, each carrying its own objective $O_i^{(i)}$ inducing a preference ordering $R_i \in L$ over composite-level alternatives A (the joint outcomes reachable through the composite’s strategic-composition machinery). The ordinal preference R_i is the ordering induced by sub-agent i ’s value functional $V_{O_i^{(i)}}(\tau)$ on the alternative set A .
- A “social welfare function” $F : L^n \rightarrow L$ is a *meta-aggregation mechanism* the designer imposes — a rule that takes the sub-agents’ ordinal preferences (or reports thereof) and produces a single social ordering on A . The designer’s task is to choose F .
- The IIA condition is the load-bearing one structurally: it requires that the social comparison of a, b depend *only* on individuals’ ordinal comparisons of a, b , not on their comparisons involving any other alternative — an *ordinal-only* informational restriction.

The Arrow theorem then states: under U + P + IIA + ND, no such F exists. The designer faces a structural impossibility on what can be aggregated over the composite of sub-agents under ordinal-only information.

Boundary characterization.

Table AQ.1

<i>Escape</i>	<i>Mechanism-design / social-choice route</i>	<i>AAT machinery?</i>
Preference-domain restriction	Single-peaked preferences (Black 1948) admit Condorcet-consistent rules including the median voter; single-crossing extends this. Domain restriction circumvents U.	Adjacent to AAT machinery. Sub-scope α' of #deriv-strategic-composition (potential / monotone game structure) is <i>adjacent to</i> single-peaked / single-crossing — the same adjacency #deriv-strategy-proofness-impossibility names. <i>Not identical</i> : sub-scope α' secures best-response convergence in the composite-dynamics, not Condorcet-consistent aggregation in the welfare function.

<i>Escape</i>	<i>Mechanism-design / social-choice route</i>	<i>AAT machinery?</i>
Weakened IIA	Positional methods — Borda count (1781), range voting, Coombs method. Each uses ranks-and-positions across the full ordering, not just pairwise comparisons.	Outside AAT machinery. Positional aggregation is not in AAT canon.
Supermajority requirements	Replace simple-majority with k/n thresholds for $k > n/2$; admits constitutional / veto-respecting aggregators (e.g., the unanimity rule, two-thirds-majority rules).	Outside AAT machinery. Supermajority-thresholding mechanisms not in AAT canon.
Cardinal preferences	Relax ordinal-only IIA; allow social aggregation over cardinal utilities — utilitarian summation, weighted utilitarianism, range voting (treating ratings as cardinal). Sen 1970, <i>Collective Choice and Social Welfare</i> , Holden-Day, §3 gives the canonical treatment of cardinal-extension escapes.	Adjacent to AAT machinery. $\#form-objective-functional$'s value-functional $V_{O_t}(\tau)$ carries cardinal value content — the value functional is real-valued, not ordinal-only. <i>Not identical</i> : AAT's value-functional formalism is per-sub-agent and per-trajectory, not a designer-side aggregation rule across sub-agents; lifting to aggregator-form requires a representational choice (utilitarian summation, weighted utilitarianism, etc.) that AAT does not currently make canonical.

Strengthened consequence. Arrow is the cluster instance with *two* AAT-side adjacencies:

Adjacency 1: sub-scope α' $\leftrightarrow$ preference-domain restriction. Shared with $\#deriv-strategy-proofness-impossibility$. Sub-scope α' of $\#deriv-strategic-composition$ (potential / monotone game structure) is adjacent to the single-peaked / single-crossing preference-domain restriction that admits Condorcet-consistent rules. The adjacency-without-identity is the same as in GS: sub-scope α' guarantees best-response *convergence* under fixed preferences, not Condorcet-consistent *aggregation* in the welfare function. The two are structurally related (both restrict preference structure to admit aggregation/convergence guarantees on a one-dimensional issue space) but operate at different layers of the composite-agent dynamics.

Adjacency 2: cardinal preferences $\leftrightarrow$ $\#form-objective-functional$. This adjacency is *unique to Arrow* among the three charter instances — GS's strategy-proofness regime is orthogonal to the ordinal/cardinal distinction (the manipulation no-go applies to any social-choice function on preference orderings regardless of whether utilities are ordinal or cardinal); MS's bilateral-trade setting already operates over cardinal valuations v_b, v_s so the Arrow IIA escape doesn't apply. For Arrow specifically, the ordinal-only IIA restriction is what AAT's value-functional formulation could *in principle* relax — the value functional V_{O_t} is real-valued and carries cardinal information per sub-agent. The adjacency is *open*: lifting AAT's per-sub-agent value-functional to a designer-side aggregator (utilitarian summation, weighted utilitarianism, range voting) requires a representational choice AAT does not currently make canonical. The open avenue is whether AAT extends to a designer-side welfare-aggregation formalism via the value-functional's cardinal content; if it does, Arrow's cardinal-preference escape would gain AAT-machinery internal content.

The other two escapes (weakened IIA via positional methods; supermajority requirements) belong entirely to social-choice theory and are honestly documented as outside AAT canon.

Epistemic Status.

Exact for the imported theorem (Arrow 1951, 1963 — classical theorem with closed-form proof in the source literature; multiple modern proofs available, e.g., Geanakoplos 2005, Reny 2001; primary-source verification scheduled for promotion past draft).

Conditional for the AAT-side translation — the composite-aggregation-mechanism reading is the natural composite-agent translation; alternative readings (treating sub-agents as factor-coalitions in a richer game; treating the social ordering as a Σ -level commitment device) exist and would map Arrow into different AAT-side vocabulary.

Robust qualitative for the boundary characterization — two named adjacencies (sub-scope $\alpha' \leftrightarrow$ preference-domain-restriction shared with GS; #form-objective-functional $\leftrightarrow$ cardinal-preferences unique to Arrow); two outside AAT canon (positional methods, supermajority).

Max attainable: *Exact* for the imported theorem. *Derived conditional* for the AAT translation. *Robust qualitative* for the boundary characterization. The cardinal-preference adjacency could lift to *derived conditional* if AAT canonicalizes a designer-side welfare-aggregator construction over the per-sub-agent value functionals; that extension is open and would be its own framework-scope decision.

Discussion.

Why this is not a #disc-identifiability-floor instance. Same actor-positioning argument as GS and MS: the actor frustrated by the Arrow no-go is the *designer* (cannot construct a welfare function satisfying all four conditions), not an agent inferring from data. The escapes are *design-constraint relaxations* (relax U via domain restriction; relax IIA via positional methods or cardinal preferences; relax majority via supermajority), not information augmentations of an inferential regime. The cluster belongs in #disc-implementation-impossibility, not in the identifiability-floor.

The cardinal-preference adjacency is structurally unique to Arrow. Among the three charter instances of #disc-implementation-impossibility, only Arrow's IIA constraint produces a cardinal/ordinal distinction that AAT's value-functional formalism could engage. GS's strategy-proofness no-go applies regardless of utility cardinality; MS's bilateral-trade setting already operates over cardinal valuations. Arrow alone forbids aggregation under *ordinal-only* information, which is the constraint AAT's real-valued V_{O_t} could relax. The adjacency is *open* — AAT does not currently formalize a designer-side aggregator that consumes the per-sub-agent value functionals — but it is the open avenue most specifically illuminating about how the framework might engage social-choice theory in the future.

Peer voice with the social-choice literature. The imported theorem is classical; the proof structure (Arrow 1951 via the decisive-set / dictator-pivotal-position argument; Geanakoplos 2005's three brief proofs; Reny 2001's geometric/topological unification with GS) is mature. AAT does not re-prove the theorem; AAT does not displace the standard escape routes; AAT does not claim that the value-functional formalism solves social-choice theory. The segment's contribution is honest recognition of the two adjacencies — sub-scope α' (shared with GS) and #form-objective-functional's cardinal value content (unique to Arrow) — with the imported theorem cited under its original name and full bibliographic references.

Relationship to #form-objective-functional. The cardinal-preference adjacency raises an open question about the value-functional's downstream extension. #form-objective-functional formalizes $V_{O_t}(\tau)$ as the agent's value functional over trajectories; the cardinal content is *within-agent* (used for the agent's own choice). Arrow's cardinal-preference escape would lift this to *cross-agent* (used by a designer to aggregate across sub-agents). Whether such a lift is a natural extension of AAT or requires new framework machinery (a designer-side aggregator separate from the per-sub-agent value functional) is the open question. #deriv-reward-channel-learning-no-go (sister to #deriv-self-actuation-grounding) recognizes the *single-interface narrowness* of V_{O_t} at the within-agent layer; the parallel cross-agent narrowness question (whether V_{O_t} aggregates honestly across sub-agents) is structurally adjacent but not yet posed in canon.

Findings.

The Arrow Translation and Its Two AAT-Side Adjacencies. **Brief:** Arrow’s theorem says that no social welfare function — no rule for aggregating individuals’ preferences into a single social preference ordering — can simultaneously satisfy four reasonable demands: be defined for every possible profile (unrestricted domain), respect unanimous preferences (Pareto efficiency), depend on each pair only through how individuals rank that pair (independence of irrelevant alternatives), and not be dictatorial (no single individual always decides). The classical escapes have been studied for decades — restrict the preference domain (single-peaked preferences admit the median voter / Condorcet rules), use a positional method (Borda count, range voting), require supermajorities, or allow cardinal preferences (weighted utilitarianism rather than ordinal-only aggregation). The contribution here is two named adjacencies on the AAT side: the preference-domain-restriction escape is *adjacent to* sub-scope α' of the strategic-composition machinery — same adjacency as GS — and the cardinal-preference escape is *adjacent to* the framework’s value-functional formalism (which carries cardinal value content per sub-agent). The cardinal adjacency is structurally unique to Arrow: GS’s strategy-proofness is orthogonal to the cardinal/ordinal distinction; MS already operates over cardinal valuations. Both adjacencies are *adjacent without being identical* — open avenues if the framework extends in those directions, not current elevations of AAT machinery to load-bearing. The other two escapes (positional methods, supermajority) belong entirely to social-choice theory and are honestly documented as outside AAT canon. The map of where the framework meets Arrow is the contribution; the theorem is classical.

Impact: Closes the third charter-instance slot of #disc-implementation-impossibility with the cluster’s *richest* AAT-side adjacency map — two named adjacencies, one shared with GS (sub-scope $\alpha' \leftrightarrow$ preference-domain restriction) and one unique to Arrow (#form-objective-functional $\leftrightarrow$ cardinal preferences). The Arrow-unique cardinal-preference adjacency is the open avenue most specifically illuminating about how the framework might engage social-choice theory in the future: AAT’s per-sub-agent value functional V_{O_t} carries cardinal value content, and lifting it to a designer-side cross-agent aggregator (utilitarian summation, weighted utilitarianism) is the natural extension that would give the Arrow cardinal-preference escape AAT-machinery internal content. Whether such an extension is a natural growth of AAT or requires new framework machinery is the open question raised by the adjacency. Together with the GS instance (named adjacency, no open avenue) and MS instance (no current adjacency, two open avenues), the cluster’s three charter instances span the full range of how AAT meets the designer-side impossibility cluster.

Novelty Claim: *Claim recognition* of the AAT-side translation of Arrow’s theorem under the composite-aggregation-mechanism reading (sub-agents per #scope-composite-agent; preference orderings induced by per-sub-agent $V_{O_t(i)}$; designer-chosen aggregator F producing social ordering); *claim differentiation* on the cardinal-preference adjacency being structurally unique to Arrow within the cluster — GS’s strategy-proofness is orthogonal to the cardinal/ordinal distinction; MS’s bilateral-trade setting already operates over cardinal valuations; only Arrow’s ordinal-only IIA constraint produces a cardinal/ordinal distinction that AAT’s value-functional formalism could engage. The imported theorem is classical (Arrow 1951, 1963); the AAT-side translation, the two-adjacency boundary characterization, and the cardinal-preference adjacency as the Arrow-unique open avenue are the contributions.

Related Work:

Table AQ.2

<i>ASF concern</i>	<i>Prior-art language</i>	<i>Relationship / Positioning</i>
The imported theorem	Arrow, K. J. (1951, 2nd ed. 1963), <i>Social Choice and Individual Values</i> , Yale University Press (found pre-2026)	<i>formal antecedent</i> — primary-source verification scheduled for promotion past draft

ASF concern	Prior-art language	Relationship / Positioning
Modern proof structure	Geanakoplos, J. (2005), “Three brief proofs of Arrow’s impossibility theorem,” <i>Econ. Theory</i> 26:211–215 (found pre-2026); Reny, P. J. (2001), “Arrow’s theorem and the Gibbard-Satterthwaite theorem: a unified approach,” <i>Econ. Letters</i> 70:99–105 (found 2026-05-20)	<i>adjacent literature</i> — Geanakoplos for the constraint-regime formulation; Reny for the geometric-topological unification with GS that informs #disc-implementation-impossibility’s mechanism-family discussion
Single-peaked preferences as the domain-restriction escape	Black, D. (1948), “On the rationale of group decision-making,” <i>J. Polit. Econ.</i> 56:23–34 (found pre-2026)	<i>adjacent literature</i> — the canonical escape; the sub-scope α' adjacency is named relative to this (shared with #deriv-strategy-proofness-impossibility)
Cardinal-preference escapes via weighted utilitarianism	Sen, A. (1970), <i>Collective Choice and Social Welfare</i> , Holden-Day, §3 (found pre-2026)	<i>adjacent literature</i> — the canonical treatment of cardinal-extension escapes; the #form-objective-functional adjacency is named relative to this (unique to Arrow within the cluster)
Borda count (positional method escape)	Borda, J.-C. (1781), “Mémoire sur les élections au scrutin,” <i>Histoire de l’Académie Royale des Sciences</i> (found pre-2026)	<i>adjacent literature</i> — outside AAT machinery; documented as honest scope-mark

Search Log:

- 2026-05-22 (*intuition-only* on the cardinal-preference adjacency): The cardinal-preference escape’s connection to #form-objective-functional is named here for the first time; the natural correspondence (AAT’s real-valued V_{O_t} carries cardinal value content per sub-agent; lifting to a designer-side aggregator is a representational choice) is intuitive but not derived. Targeted future search candidates: utilitarian and weighted-utilitarian social welfare functions (Harsanyi 1955, *J. Polit. Econ.* 63:309–321; Sen 1970 §3 cardinal-comparability spectrum); range voting and ratings-based methods (Smith 2000 and follow-ups); whether the per-sub-agent value functional’s cardinal content can be made canonically aggregable without introducing new framework machinery.
- 2026-05-22 (*intuition-only* on the composite-aggregation-mechanism reading): Same composite-revelation-mechanism reading as #deriv-strategy-proofness-impossibility. Alternative readings exist (treating F as a Σ -level commitment device; treating sub-agents as factor-coalitions in a richer cooperative game); not pursued here.

Working Notes.

- **Provenance.** Authored 2026-05-22 as commit 4 of the post-Track-CR Track-B cycle, executing spikes/implementation-impossibility-meta-segment-plan.md §4.c. The 2026-05-20 strengthen-first arm record at spikes/.integrated/spike-4th-identifiability-floor-instance-2026-05-20.md §5 flagged Arrow as a candidate instance of the sister meta-segment route.
- **Open before candidate.** (a) Primary-source verification of Arrow 1951 / 1963 — the theorem statement above is taken from the textbook syntheses and modern proof papers (Geanakoplos 2005, Reny 2001); direct primary-source reading is queued. (b) The cardinal-preference $\leftrightarrow$ #form-objective-functional adjacency is currently *robust qualitative* — there is no derived identification of AAT’s value-functional with a designer-side aggregator; lifting the adjacency to *derived conditional* would require either canonicalizing a designer-side welfare-aggregation construction over the per-sub-agent $V_{O_t(i)}$, or honestly determining that no such construction is natural within AAT’s current scope. Indexed at spikes/PROPOSED.md Tier 2 (added 2026-05-24) $\rightarrow$ spikes/PROPOSED-ADVANCED.md §Phase-3 #18 (two-branch spike: constructive lift first per strengthen-first; no-go branch as honest fallback). (c) The shared sub-scope- α' adjacency with #deriv-strategy-proofness-impossibility inherits the open-derivation flagged there: whether single-peaked preferences on an issue-space *derivably* induce a potential-game structure under best-response dynamics is the parameter-level identification question; if it lands, both GS and Arrow’s sub-

scope- α' adjacencies lift to *exact* simultaneously. Indexed at spikes/PROPOSED.md Tier 3 (added 2026-05-24) under #deriv-strategy-proofness-impossibility (the canonical home for the derivation; symmetric lift here).

- **The cluster's range, complete after this commit.** The three charter instances of #disc-implementation-impossibility now span the full range of AAT/mechanism-design boundary types: GS (one named adjacency, no open avenue), MS (no current adjacency, two open avenues), Arrow (one shared named adjacency, one Arrow-unique open avenue). That range is the cluster's epistemic content — the framework recognizes where its machinery contributes, where it sits adjacent, where it has open avenues, and where the boundary is outside entirely. The cluster's discipline filter admits all three; future candidate instances (Roberts 1979 affine-maximizers, Skalse 2022 reward-hacking) would be evaluated against the same filter in a future cycle.
 - **The Arrow-unique cardinal adjacency as the cluster's most-illuminating open question.** Among all the adjacencies and avenues across the three charter instances, the cardinal-preference $\leftrightarrow$ #form-objective-functional adjacency is the one most directly illuminating about how AAT might grow into mechanism design proper. The value functional already carries cardinal content per sub-agent; the only remaining structural question is whether cross-agent aggregation is a natural extension of the framework or requires distinct designer-side machinery. Flagging this here so future framework-scope decisions can engage the question with full provenance.
-

Appendix AR

Derivation: Fisher-Local Update Gain

Under the **Fisher-local invariance regime** — smooth log-likelihood admitting non-degenerate local quadratic expansion at the current model parameter, with single-step update at first order in the step size — the natural-gradient Bayesian posterior mean shift has the form $\Delta\theta = K \cdot \tilde{\nabla}$ with **gain operator** $K = (H_M + H_L)^{-1}H_L$ and **scalar collapse** $\eta^* = U_M/(U_M + U_o)$ along the natural-gradient direction in the commuting / shared-eigenbasis case (always in 1-D; under (PI)/Čencov in the natural-gradient direction in higher dimensions). The result derives the form #emp-update-gain carries as an empirical claim, places it in the Fisher-local regime as exact, and recovers linear-Gaussian (Kalman) and conjugate-Bayesian instances as cases where the local quadratic expansion is globally exact. Companion at the model-parameter-update layer to #deriv-fisher-whitened-update-rule’s edge-update derivation: that segment derives update *direction* under correlated evidence; this segment derives update *magnitude* under the Fisher-local regime. Both share the (PI) parameterization-invariance axiom from #scope-agent-identity and Čencov 1982 uniqueness as AAT-internal forcing for the Fisher metric.

The same scalar $\eta^* = U_M/(U_M + U_o)$ falls out of three independent structural derivations — local-Gaussian Laplace expansion, Bregman / KL projection onto the local tangent plane (Pythagorean projection in the Fisher metric), and Cramér-Rao / inverse-Fisher precision-additive composition. The agreement is not coincidence: the Fisher-local regime *is* the regime where prior and likelihood compose as a Pythagorean projection in the Fisher metric. The matrix-valued K is the natural object — its eigenvalues per coordinate are $\eta_i^* = u_{M,i}/(u_{M,i} + u_{o,i})$ — and it serves as the **per-coordinate primitive** for the tensor adaptive tempo of 3.8 (scalar tempo recovered in the shared-eigenbasis collapse; tensor form preserves per-direction information the scalar averages away). The downstream consumer at the persistence layer is #deriv-matrix-persistence-condition, where cross-dimensional correction makes the matrix-Loewner condition strictly sharper than per-coordinate.

Formal Expression.

Setup — Fisher-local quadratic regime.

Let $\theta \in \mathbb{R}^d$ parameterize the agent’s model and θ_t be the current point estimate (mode of the prior π_0 , or center of expansion). Three regime conditions:

Definition (Fisher-local regime, conditions (R1)–(R3)).

- **(R1) Smooth log-likelihood admitting non-degenerate local quadratic expansion.** $\log \pi_0$ and $\log p(o \mid \cdot)$ are C^3 in a neighborhood of θ_t , with prior precision $H_M := -\nabla^2 \log \pi_0(\theta_t)$ and observed information $H_L := -\nabla^2 \log p(o \mid \theta)|_{\theta_t}$ satisfying $H_M + H_L > 0$ (joint positive-definiteness; see §”Boundary admissibility” for why this is weaker than $H_M > 0$ and $H_L > 0$ both).

- **(R2) First-order-in-step-size regime.** The posterior update $\Delta\theta = \theta_{t+1} - \theta_t$ is small enough that $O(\|\Delta\theta\|^3)$ terms in the quadratic expansion are negligible compared to the quadratic terms.
- **(R3) Bayesian-coherent update.** θ_{t+1} is taken to be a coordinate of the posterior $p(\theta | o) \propto \pi_0(\theta) \cdot p(o | \theta)$ — mean, mode, or natural-parameter (they coincide for Gaussian posteriors, which is what the quadratic expansion produces).

Under (R1)–(R3), the local log-posterior is

$$\log p(\theta | o) = \text{const} + s^T(\theta - \theta_t) - \frac{1}{2}(\theta - \theta_t)^T(H_M + H_L)(\theta - \theta_t) + O(\|\theta - \theta_t\|^3), \quad (\text{AR.1})$$

where $s := \nabla_\theta \log p(o | \theta)|_{\theta_t}$ is the score. The posterior is approximately $\mathcal{N}(\theta^\star, (H_M + H_L)^{-1})$ with mean shift

Derived (posterior-mean-shift).

$$\theta^\star - \theta_t = (H_M + H_L)^{-1}s. \quad (\text{AR.2})$$

The AAT-vocabulary correspondences are

$$U_M := H_M^{-1}, \quad U_o := H_L^{-1} \quad (\text{AR.3})$$

— model uncertainty is the inverse prior precision, observation uncertainty is the inverse Fisher / inverse observed information (Cramér-Rao floor on what the observation pins down). The matrix forms generalize #emp-update-gain's scalar U_M, U_o in the natural way.

Gain decomposition on the natural-gradient direction.

The **natural gradient** of the log-likelihood at θ_t in the Fisher metric is

Definition (natural gradient at the current point).

$$\tilde{\nabla} \log p(o | \theta_t) := \mathcal{I}(\theta_t)^{-1}s. \quad (\text{AR.4})$$

For a single observation in the Fisher-local regime, $\mathcal{I}(\theta_t) = H_L$ up to the observed-vs-expected information distinction (standard under (R1)+regularity); the derivations below use H_L directly as the Fisher.

Rewriting the posterior mean shift in the natural-gradient direction:

Derived (gain-decomposition).

$$\Delta\theta = (H_M + H_L)^{-1}s = (H_M + H_L)^{-1}H_L \cdot (H_L^{-1}s) = K \cdot \tilde{\nabla}, \quad (\text{AR.5})$$

with **gain operator**

$$K := (H_M + H_L)^{-1}H_L. \quad (\text{AR.6})$$

The matrix-valued K is the natural object. The scalar collapse follows in the commuting / shared-eigenbasis case (always in 1-D; under (PI)/Čencov along the natural-gradient direction in higher dimensions): eigenvalues of K are

$$\eta_i^* = \frac{h_{L,i}}{h_{M,i} + h_{L,i}} = \frac{1/u_{o,i}}{1/u_{M,i} + 1/u_{o,i}} = \frac{u_{M,i}}{u_{M,i} + u_{o,i}} \quad (\text{AR.7})$$

per coordinate, collapsing to the scalar

$$\boxed{\eta^* = \frac{U_M}{U_M + U_o}} \quad (\text{AR.8})$$

— the form `#emp-update-gain` carries, derived from the Fisher-local quadratic expansion of prior and likelihood and applied to the natural-gradient direction.

Boundary admissibility — improper-prior and degenerate-likelihood directions.

The gain operator $K = (H_M + H_L)^{-1}H_L$ is well-defined whenever $H_M + H_L > 0$ — not necessarily when both H_M and H_L are individually positive-definite. This admits **degenerate-likelihood directions** (observation uninformative along some coordinates, H_L rank-deficient): K 's eigenvalue along those directions is 0, $u_{o,i} \rightarrow \infty$, $\eta_i^* \rightarrow 0$, no update along that direction — *trust the prior where the observation says nothing*. It also admits **improper-prior directions** (degenerate prior, $H_M = 0$ along some coordinate) when H_L pins the direction: $\eta_i^* \rightarrow 1$ along uninformative-prior directions where the observation is informative — *trust the observation where the prior says nothing*, the maximum-likelihood limit.

Derived (admissibility-of-degenerate-curvature).

The minimal regime condition is $H_M + H_L > 0$, with $U_M/(U_M + U_o)$ defined per-eigendirection by the appropriate limit. This corresponds to standard improper-prior admissibility in Bayesian inference and to standard observation-non-identifiability admissibility in maximum-likelihood estimation.

Recovery of linear-Gaussian and conjugate-Bayesian cases.

When the quadratic expansion in (R1) holds *globally* — not just locally — the derivation gives η^* without any expansion error. Two canonical cases:

Derived (global-exactness-special-cases).

- **Linear-Gaussian (Kalman).** Prior $\mathcal{N}(\theta_i, U_M)$, likelihood $\mathcal{N}(o; H\theta, U_o)$ with linear observation operator H : the log-posterior is exactly quadratic. The scalar Kalman gain $K = U_M H^T (H U_M H^T + U_o)^{-1}$ collapses to $\eta^* = U_M/(U_M + U_o)$ in the scalar case (Kalman 1960; Bishop 2006 §3.3). Exact globally.
- **Conjugate Bayesian (Beta-Bernoulli, Dirichlet-multinomial, Gaussian-Gaussian, exponential-family natural-parameter conjugacy).** Natural-gradient VI is the exact posterior update at step size 1 for conjugate families (Khan-Lin 2017 conjugate-computation variational inference). The Fisher-local expansion is globally exact in natural-parameter coordinates; $\eta^* = U_M/(U_M + U_o)$ holds exactly with U_M, U_o read from the conjugate sufficient-statistic precision counts.

For general smooth models, the natural-gradient invariance theorem of Amari 1998 guarantees that the natural gradient is parameterization-invariant; the Fisher-local quadratic expansion is the order at which the natural gradient and the exact posterior agree to first order in step size.

Epistemic Status.

Conditional. Max attainable: *exact* under (R1)–(R3) + (PI) parameterization-invariance from #scope-agent-identity. Under (PI) and Čencov 1982 uniqueness, the Fisher metric is the AAT-internally canonical metric on the statistical manifold (the 4th primary instance of #disc-additive-coordinate-forcing); the natural-gradient direction is therefore the AAT-internally canonical reference direction along which η^* reads as $U_M/(U_M + U_o)$. The textbook information-geometry result (Amari 1998 natural-gradient invariance theorem) is imported; the (PI)/Čencov forcing in #deriv-fisher-whitened-update-rule’s Path A makes the import AAT-internal.

Linear-Gaussian (Kalman) and conjugate-Bayesian instances are cases where the local quadratic expansion is *globally* exact, so $\eta^* = U_M/(U_M + U_o)$ holds without truncation error.

Outside the Fisher-local regime — heavy-tailed posteriors (no second moment), structurally non-smooth likelihoods, multimodal uncertainty where the local quadratic misrepresents global structure — the quantitative form is *not* derived here. The *qualitative* direction (gain rises with model uncertainty, falls with observation noise) is preserved as *robust qualitative* from #emp-update-gain’s broader empirical scope; the first-order form is recovered locally; what need not hold is global quantitative fidelity.

This segment depends on (PI) from #scope-agent-identity only. The triple (PI)+(R)+(K) of the Markov-morphism layer (carried by the Fisher-Rao bias-bound machinery; the (R) Riemannian-structure and (K) KL-second-order-matching axioms are stronger) is *not* invoked here. If (R) and (K) are asserted at the scope-level in a future cycle, this segment’s status remains conditional on (PI) alone unless explicitly extended.

Discussion.

Why read η^* off the natural-gradient direction. The Bayesian-coherent posterior mean shift $\Delta\theta = (H_M + H_L)^{-1}s$ is a fixed vector; its decomposition into “gain times reference direction” depends on the choice of reference direction. Two natural choices give two faces of the same shift:

- **Natural-gradient direction** $\tilde{V} = H_L^{-1}s$. Coefficient: $\eta_{\text{NG}}^* = H_L/(H_M + H_L) = U_M/(U_M + U_o)$. *AAT-internally canonical* under (PI)/Čencov — the Fisher metric is the unique Markov-invariant metric, so the natural gradient is the unique coordinate-invariant gradient direction (Čencov 1982; extended by Ay-Jost-Lê-Schwachhöfer 2017).
- **Prior-curvature-rescaled direction** $H_M^{-1}s$. Coefficient: $\eta_{\text{prior}}^* = H_M/(H_M + H_L) = U_o/(U_M + U_o)$. The *complement*; the same algebraic decomposition arises in Kalman literature as the prior weight in the convex combination “posterior = (1-K) prior + K observation.”

The two coefficients are duals on the same posterior shift; their product reduces to the posterior covariance $(H_M + H_L)^{-1}$. AAT names the natural-gradient coefficient as the canonical η^* because (PI)/Čencov picks out the natural-gradient direction as the unique coordinate-invariant choice.

Three derivation routes converge. The same scalar $\eta^* = U_M/(U_M + U_o)$ falls out of three independent structural derivations:

Table AR.1

Route	Lever	Surface meaning of η^*
Local-Gaussian Laplace expansion (§Setup–Gain decomposition above)	Algebraic completion of squares on additive log-posterior quadratics	Posterior mean shift coefficient on the natural-gradient direction

<i>Route</i>	<i>Lever</i>	<i>Surface meaning of η^*</i>
Bregman / KL projection onto the local tangent plane	First-order condition on the variational free energy along natural-gradient tilt	Projection coefficient of the posterior onto the local tangent plane (Pythagorean projection in Fisher metric per Amari-Nagaoka 2000 §3.2)
Cramér-Rao / inverse-Fisher	Precision-additive composition $1/U_M + 1/U_O$, inverted	Fractional observation contribution to posterior precision

The agreement is not coincidence: the Fisher-local regime *is* the regime where prior and likelihood compose as a Pythagorean projection in the Fisher metric — the same exponential-family / Bregman / Pythagorean structure underlying #deriv-strategy-cost-regret-bound's ρ -decomposition. The three routes are three faces of one geometric object: precision-additive composition of two Gaussian information sources, read off the natural-gradient direction.

Sibling positioning vs #deriv-fisher-whitened-update-rule. That segment lives at the **edge-update layer** of Part II's strategy DAG: it derives the Fisher-whitening correction for the *direction* of edge updates under correlated evidence (L1'/L2). This segment lives at the **model-parameter-update layer** of Part I: it derives the *magnitude* of the natural-gradient Bayesian update via the Fisher-local invariance regime. Both depend on (PI) + Čencov as the AAT-internal axiom forcing the Fisher metric. Together they make the Fisher-local invariance regime AAT-internally complete: direction (whitening) at the edge layer and magnitude (gain) at the model-parameter layer are both derived from the same axiomatic chain.

Special case of #deriv-adaptive-gain-dynamics' meta-gain framework. The gain operator $K = (H_M + H_L)^{-1}H_L$ is a *deterministic function of the current state* — it depends on the prior covariance H_M^{-1} (part of the agent's M_t state) and on the observation Fisher H_L (depends on θ_t and o). It is therefore the **deterministic-meta-gain** special case of #deriv-adaptive-gain-dynamics' (MG-1)–(MG-4) machinery: the meta-gain is determined by the primary state rather than independently learned. All four meta-gain conditions are satisfied trivially (symmetric-positive-definite K on the interior; smoothness for exponential families; bounded primary-meta coupling). Resolving the epistemic-opacity question that #def-observation-function axiomatizes — i.e., the agent does not know U_o a priori and must estimate it from its own mismatch sequence — is one level up: it adds the Mehra-style meta-channel that #deriv-adaptive-gain-dynamics treats as its primary instance.

Downstream tempo and persistence machinery. #emp-update-gain's Discussion §"Connection to adaptive tempo" names $\mathcal{T} = \nu \cdot \eta^*$. The derivation here promotes η^* 's tier from empirical / robust-qualitative (over the cross-domain validity tail — RL, PID, software-developer) to derived (conditional on Fisher-local regime) for the Kalman / conjugate / natural-gradient core. The tempo product $\mathcal{T} = \nu \cdot \eta^*$ inherits the regime distinction directly; the persistence-condition machinery (#result-persistence-condition, #result-sector-condition-stability) and the adaptive-tempo / aporia diagnostics (#def-adaptive-tempo, #def-aporia) all gain the tighter η^* exactness statement in their Fisher-local instances. The qualitative direction claim and the failure-mode framing (gain-collapse $\rightarrow$ epistrophe failure) survive unchanged outside the Fisher-local regime.

Tensor adaptive tempo connection. The matrix gain operator K is the per-coordinate primitive for tensor-valued adaptive tempo: $\mathcal{T} = \nu \cdot K$ as a matrix product, with the existing scalar $\mathcal{T} = \nu \cdot \eta^*$ recovered in the shared-eigenbasis limit. See #def-adaptive-tempo's Tensor extension.

Working Notes.

- **Multi-step extrapolation rate.** The Fisher-local exactness is a one-step statement. Iterating over many observations accumulates higher-order curvature error in the non-Gaussian / non-conjugate case. Bernstein-von Mises asymptotics (van der Vaart 1998 §10.2) give a $1/\sqrt{n}$ posterior concentration rate; whether η^* 's exactness inherits a matching $1/\sqrt{n}$ degradation outside Kalman / conjugate is an open quantitative-rate question. Open spike candidate.

- **Edgeworth higher-order correction.** The $O(\|\Delta\theta\|^3)$ correction to η^* in terms of third / fourth derivatives of the log-likelihood admits a standard Edgeworth-expansion treatment. Worth a follow-up spike if it becomes load-bearing for any AAT use case.
- **Variational-bound extension.** Under variational approximation $\text{KL}(q\|p) \leq \epsilon$, the gain inherits a Pinsker-tight degradation factor analogous to #deriv-variational-sector-condition's $O(\sqrt{\epsilon})$ B1 degradation. Composing this with the derivation gives η^* in sub-scope α' under approximate-posterior agents.
- **Multimodal posteriors.** When the prior or posterior is multimodal, the local-quadratic expansion around any one mode misses the global structure. Two natural moves: per-mode $U_M/(U_M + U_o)$ with mode-mixture weights as a separate state; or accept the qualitative direction claim only. Worth a follow-up spike if multimodal posteriors become load-bearing.
- **Consolidation / replay connection.** #form-consolidation-dynamics-style between-event consolidation drives U_M up (model uncertainty grows under offline replay via IB-gap reduction). The next online observation gets up-weighted because $U_M/(U_M + U_o)$ rises with U_M . This is the formal connection between consolidation-induced uncertainty growth and increased online responsiveness — worth tightening into a result if #form-consolidation-dynamics lands a quantitative form.
- **Honest obstructions.** (O1) The natural-gradient invariance theorem (Amari 1998) is imported, not derived; the (PI)/Čencov forcing in #deriv-fisher-whitened-update-rule's Path A makes the import AAT-internal. Without (PI), the natural-gradient direction is a chosen-not-forced direction and the derivation degrades to “the form on the chosen direction.” (O2) Single-observation per-step setup; multi-observation batches generalize mechanically ($H_L = \sum_i H_{L,i}$ for independent observations). (O3) Matrix vs scalar — the matrix K is the natural object; the scalar $U_M/(U_M + U_o)$ is its eigenvalue in the commuting / 1-D / shared-eigenbasis case. (O4) Step-size boundary qualitative; the actual Fisher-local cutoff depends on the third-derivative norms of $\log \pi_0$ and $\log p(o \mid \cdot)$. (O5) Robust-qualitative downstream survives outside the Fisher-local regime by the general structure of Bayesian updating, not by this derivation.
- **Cross-reference to Paper 3 chart-rescaling no-go.** The (PI) dependence is forced by the chart-rescaling no-go on Euclidean chart norms (NeurIPS 2026 Paper 3, “How Much Can LLMs Hallucinate?”, §4 Theorem 4.2 / #thm-no-go); outside (PI), no universal-constant claim survives. The natural-gradient direction inherits this forcing; the canonical-direction argument depends on (PI), not on (R) or (K) of the Markov-morphism triple. See spikes/neurips-back-integration-2026-05-08.md §1 Paper 3 entry #3 and the source ~/src/neurips/03-llm-hallucinate-bound/.
- **Landing context.** Landed in the 2026-05-12 audit-strengthening cycle (AAT-5); see CHANGELOG 2026-05-12. The load-bearing three-route convergence and boundary-admissibility content is in the Derivation and Discussion above; the originating spike is absorbed archaeology, not a live reference.
- **Downstream consumer at the persistence layer.** The matrix gain operator $K = (H_M + H_L)^{-1}H_L$ derived here is the per-channel primitive that #def-adaptive-tempo's Tensor extension aggregates into $\mathcal{T} = \sum_k \nu^{(k)}K^{(k)}$. The persistence-layer consumer is the matrix-Loewner persistence condition in #deriv-matrix-persistence-condition: under matrix $\mathcal{T}$ and matrix disturbance Σ_w , the agent persists iff $\mathcal{T}$ is Hurwitz and the stationary covariance Σ_∞ (solving $\mathcal{T}\Sigma_\infty + \Sigma_\infty\mathcal{T}^T = \Sigma_w$) is strictly Loewner-dominated by $D_\delta = \text{diag}(\delta_{\text{critical},k}^2)$. The matrix-Loewner form is *strictly sharper* than per-coordinate: under cross-dimensional correction (off-diagonal $\mathcal{T}$ — the generic Fisher-local case when prior and likelihood do not share the coordinate basis), per-coordinate gives a false-pass while matrix-Loewner reads persistence correctly off the worst direction.

Appendix AS

Derivation: Scope of Tempo Additivity — Redundancy, Synergy, and the Echo-Chamber Bound

The additive tempo formula $\mathcal{T} = \sum_k \nu^{(k)} \eta^{(k)*}$ (3.8) is exact when the observation channels carry cross-channel-independent noise, and its deviation from the true joint capacity under correlated noise is **signed**: correlated channels can deliver *less* joint information than the additive accounting claims (redundancy — the additive form overcounts) or *more* (synergy — correlated noise can be cancelled, and the additive form undercounts). This segment derives the exact deviation in the Fisher-local regime, proves the additive form is an exact upper bound with an explicit, closed-form redundancy penalty in the **common-source noise regime** (the “echo chamber”: a shared persistent bias contaminating several channels), proves an information-**saturation** theorem for persistent shared bias (no number of correlated channels buys information past the shared-bias floor), and establishes two no-gos: no sign-blind dependence measure (conditional mutual information, total correlation) can be the exact correction term, and no convention-free channel-local attribution of the deviation exists for three or more channels (the partial-information-decomposition obstruction). The joint tempo and the *total* deviation are canonical; per-channel corrected contributions are not.

Formal Expression.

Setup — the information reading of tempo.

Work in the Fisher-local / Gaussian location regime of [AR](#): environment parameter θ , model uncertainty U_M (posterior variance), channels $k = 1, \dots, n$ with independent Poisson event rates $\nu^{(k)}$ and per-event observations $o^{(k)} = \theta + n^{(k)}$, noise variance $U_o^{(k)} = \text{Var}(n^{(k)})$. Let $\mathfrak{F}$ denote the posterior-precision inflow rate delivered by the *optimal joint filter* (Fisher information about θ acquired per unit time). The **effective tempo** is the fractional uncertainty-contraction rate

Definition (effective tempo, Fisher-local regime).

$$\mathcal{T}_{\text{eff}} := -\frac{d}{dt} \log U_M(t) = U_M \cdot \mathfrak{F}, \quad (\text{AS.1})$$

the second equality since $\frac{d}{dt} (1/U_M) = \mathfrak{F}$. Single-channel consistency: $\mathfrak{F} = \nu/U_o$ gives $\mathcal{T}_{\text{eff}} = \nu U_M/U_o$, which agrees with $\nu\eta^* = \nu U_M/(U_M + U_o)$ to first order in the per-event information ratio U_M/U_o — the same first-order-in-step-size regime (R2) under which [AR](#) operates and under which tempo is a rate at all. All statements below are exact at the Fisher-information level and transfer to gain-form tempo at that order.

Additivity under noise independence.

If the channel noises are independent across channels and events (conditional independence of the pooled event stream given θ), scores add and score cross-covariances vanish, so Fisher information is exactly additive:

Derived (additivity-under-independence).

$$\mathfrak{F}_{\text{joint}} = \sum_k \nu^{(k)} \frac{1}{U_o^{(k)}} = \mathfrak{F}_{\text{add}} \implies \mathcal{F}_{\text{eff}} = \sum_k \nu^{(k)} \eta^{(k)*} \text{ (first order).} \quad (\text{AS.2})$$

Independence is **sufficient but not necessary** for equality (see the harmonic-mean hypersurface below).

Independent Poisson streams almost surely have no simultaneous events. If noise is *fresh per event* (i.i.d. across events), instantaneous cross-channel noise correlation is never sampled and additivity remains exact regardless of it. What breaks additivity is noise dependence **across events** — a persistent shared noise component. The redundancy question is governed by the ratio of the noise-correlation timescale to event interarrival times.

Derived (asynchrony refinement).

The exact deviation, and its indefinite sign.

For a synchronized batch (one event per channel, jointly Gaussian noise Σ_n , $(\Sigma_n)_{kk} = U_o^{(k)}$), the joint Fisher information and the deviation from additivity are

Derived (signed-deviation).

$$J_{\text{joint}} = \mathbf{1}^\top \Sigma_n^{-1} \mathbf{1}, \quad \Delta := \sum_k \frac{1}{U_o^{(k)}} - J_{\text{joint}}, \quad (\text{AS.3})$$

and for two channels with noise covariance c (writing U_1, U_2 for the variances):

$$J_{\text{joint}} = \frac{U_1 + U_2 - 2c}{U_1 U_2 - c^2}, \quad \Delta = \frac{c(2U_1 U_2 - c(U_1 + U_2))}{U_1 U_2 (U_1 U_2 - c^2)}. \quad (\text{AS.4})$$

$\Delta = 0$ iff $c = 0$ or $c = 2U_1 U_2 / (U_1 + U_2)$ (the harmonic mean of the variances — admissible whenever $U_1 \neq U_2$): equality holds on a strictly larger set than independence, with redundancy and synergy cancelling exactly. Δ has **no definite sign**:

- **Anti-correlated noise (synergy).** $U_1 = U_2 = \sigma^2, \rho < 0$: $J_{\text{joint}} = 2/(\sigma^2(1 + \rho)) > 2/\sigma^2$. Averaging cancels the noise; the additive form *undercounts*.
- **Positively correlated, heterogeneous variances (synergy).** $U_1 = 1, U_2 = 100, c = 9.9$: $J_{\text{joint}} = 81.2/1.99 \approx 40.8 \gg 1.01$. The noisy channel serves as a noise *reference*; the combination $o^{(1)} - (c/U_2) o^{(2)}$ nearly cancels channel 1's noise.
- **Nonnegative equal-variance correlation, $n = 3$ (synergy).** Correlations $\rho_{12} = \rho_{23} = 0.7, \rho_{13} = 0$ (positive-definite): $\mathbf{1}^\top R^{-1} \mathbf{1} = 10 > 3$.

So “correlated channels $\Rightarrow$ the additive formula overcounts” is false as a general claim; overcounting is the *redundant* half of a two-sided phenomenon whose other half is noise-cancelling synergy.

The echo-chamber regime — redundancy penalty and saturation.

Common-source noise model: $n^{(k)} = s + w^{(k)}$ with shared bias $s \sim \mathcal{N}(0, \sigma_s^2)$ and independent $w^{(k)} \sim \mathcal{N}(0, \sigma_{w,k}^2)$, so $\Sigma_n = \sigma_s^2 \mathbf{1}\mathbf{1}^\top + \text{diag}(\sigma_{w,k}^2)$. By Sherman-Morrison, with $q_k := 1/\sigma_{w,k}^2$, $q := \sum_k q_k$, and $f(x) := x/(1 + \sigma_s^2 x)$:

Derived (common-source-redundancy-penalty).

$$J_{\text{joint}} = f(q) \leq \sum_k f(q_k) = \sum_k \frac{1}{\sigma_s^2 + \sigma_{w,k}^2} = \sum_k J^{(k)}, \quad (\text{AS.5})$$

by strict concavity of f with $f(0) = 0$ (strict subadditivity), with equality iff $\sigma_s^2 = 0$ or at most one channel is active. The **redundancy penalty** is explicit and closed-form: $R_s = \sum_k f(q_k) - f(q) > 0$. In this regime the additive tempo *is* an exact upper bound on effective tempo, for arbitrarily heterogeneous channels.

If the shared bias s is **persistent across events**, every observation measures $\theta + s$ plus fresh noise: the pooled data identify $\theta + s$ at additive rate, but the posterior precision about θ itself is bounded for all time by $1/U_M^{(0)} + 1/\sigma_s^2$, since $\text{Var}(\theta \mid \theta + s) = (1/U_M^{(0)} + 1/\sigma_s^2)^{-1}$ and conditioning on the data is coarser than conditioning on $\theta + s$. Additive accounting claims precision growing without bound; the true joint information **saturates** at the shared-bias floor. No number of common-source channels, at any event rate, buys information past $1/\sigma_s^2$.

Derived (saturation-under-persistent-bias).

Two no-gos.

No exact correction of the form $\mathcal{J}_{\text{eff}} = \mathcal{J}_{\text{add}} - g(D)$ exists for any dependence measure D invariant under noise-correlation sign flip — conditional mutual information $I(e^{(1)}; e^{(2)} \mid M_{\tau-})$, total correlation (Watanabe 1960), or any such functional. Proof: ρ and $-\rho$ in the equal-variance two-channel family give identical Gaussian conditional mutual information $-\frac{1}{2} \log(1 - \rho^2)$ but different joint informations $2/(\sigma^2(1 \pm \rho))$. The deviation depends on the dependence *geometry*, not the dependence *strength*; the exact correction is the Fisher deficit Δ , which requires Σ_n , not a scalar dependence summary. Conditional mutual information survives as a **witness**: in the Gaussian location model $\Delta \neq 0 \Rightarrow I > 0$ (converse false on the harmonic-mean hypersurface), and in the common-source regime both I and R_s increase with σ_s^2 .

Derived (no-sign-blind-penalty).

A channel-local, convention-free additive decomposition of the joint information rate into per-channel redundancy-corrected contributions is the demand of the **partial information decomposition** program, and it is known not to exist in general: Williams & Beer 2010 define the redundancy lattice; Bertschinger, Rauh, Olbrich, Jost & Ay 2014 give the two-source unique-information measure; Rauh, Bertschinger, Olbrich & Jost 2014 prove no nonnegative decomposition on the Williams-Beer lattice satisfying the identity axiom exists for three or more sources; Barrett 2015 shows the Gaussian collapse of the main two-source candidates. Consequence: $\mathcal{J}_{\text{joint}}$ and the total deviation Δ are canonical and exact; “channel k ’s corrected tempo contribution” is not a well-defined quantity for $n \geq 3$, and any AAT use of one must declare its attribution convention explicitly.

Derived (no-canonical-channel-local-decomposition — imported).

Epistemic Status.

Conditional. The derivations are exact within the Fisher-local / Gaussian location regime of AR (conditions (R1)–(R3)) plus the stated noise models (jointly Gaussian batch noise for the deviation formula; the common-source structure for the penalty and saturation theorems); the transfer from Fisher-information statements to gain-form tempo is first-order in the per-event information ratio, the same regime in which tempo is a rate. Additivity-under-independence is general (standard Fisher additivity, not merely Gaussian); the counterexamples are exact and finite, computed in-segment. The two-source PID landscape and the $n \geq 3$ impossibility are imported results, cited by their real names; the mapping of the PID obstruction onto channel-local tempo attribution is AAT’s reading and is stated at that grade. Citation details for the imported PID results are flagged for primary-source verification before promotion past draft (see Working Notes). This segment supplies the derivation behind the channel-independence condition named in 3.8’s Epistemic Status; it *refutes* the previously asserted general upper bound (the deviation is signed) and *establishes* the bound on its true scope (common-source regime), so downstream uses of “additive tempo is conservative” are unsafe unless the common-source (or general nonpositive- Δ) structure is verified for the application.

Discussion.

What the correction to the prior assertion changes. The earlier claim — additive tempo is always an upper bound, equality iff channels informationally independent — read channel correlation as pure overlap. The derivation shows correlation is two-sided: overlapping *information* (common source) makes the additive form overcount, but correlated *noise* is also a resource — a correlated reference channel lets the filter cancel noise it could never average away, and then the additive form undercounts, potentially by large factors. Both directions matter for agent design: the redundant direction is the echo-chamber warning (diversity of observation *sources*, not count of observers), and the synergistic direction says a channel that is nearly useless in isolation (huge U_0 , negligible η^*) can be extremely valuable jointly — a reason to retain “bad” sensors that are correlated with good ones’ failure modes.

The echo chamber, quantitatively. The common-source saturation theorem is the formal content of the organizational intuition: three VPs reading one broken dashboard are one channel with extra steps, and past the shared-bias floor their combined reports carry zero marginal information about the world. A filter that *models* the channels as independent additionally manufactures false confidence — its internal U_M shrinks at the additive rate while true error stalls at the bias floor, driving $\eta^* \rightarrow 0$ while genuinely wrong: the gain-collapse pathology of 3.6 arrived at through mis-modeled channel structure rather than miscalibrated uncertainty [Discussion].

Why no simple penalty formula exists. The no-gos localize the obstruction precisely. A scalar dependence summary cannot carry the correction (sign-blindness), and a per-channel attribution cannot be made canonical for $n \geq 3$ (PID impossibility). What *is* available exactly: the joint object $\mathbf{1}^\top \Sigma_n^{-1} \mathbf{1}$ (tensor case: $A^\top \Sigma_n^{-1} A$) whenever the noise covariance structure is known or estimable, and the closed-form penalty on the common-source scope. This is a scope-precision result in the CS sense: the boundary of the additive form is now characterized rather than gestured at.

Anisotropy is orthogonal. With multivariate θ and channel maps A_k , cross-channel noise independence gives exact *matrix* additivity $\mathcal{J} = \sum_k \nu^{(k)} A_k^\top (\Sigma_n^{(k)})^{-1} A_k$ — channels “measuring different aspects” is handled by the tensor form of 3.8, not by redundancy corrections. Redundancy/synergy is an off-block-diagonal-noise phenomenon; anisotropy is an eigenbasis/observation-map phenomenon; the common-source subadditivity holds in the Loewner order by the same Sherman-Morrison-Woodbury argument.

Composition consequence. 15.1 inherits both directions: allied reports drawing on one upstream source are the common-source regime (additive communication tempo overcounts, saturating at the source’s bias floor), while allies with anti-correlated errors triangulate — a genuinely super-additive composition effect the additive accounting misses in the favorable direction.

Working Notes.

- **Citation verification before promotion.** The PID citations (Williams-Beer 2010 arXiv:1004.2515; Bertschinger-Rauh-Olbrich-Jost-Ay 2014 *Entropy* 16(4); Rauh et al. 2014 ISIT; Barrett 2015 *PRE* 91; Watanabe 1960 *IBM J. Res. Dev.* 4) are from literature memory and must be primary-source-verified (relata pass) before claims-verified. The mathematical content here does not lean on them — they scope the no-go's external grounding.
 - **Finite-correlation-time interpolation.** §Asynchrony gives the two exact endpoints (fresh noise: additive exact; persistent bias: saturation). The intermediate regime (noise correlation time comparable to interarrival times) should interpolate via an effective $\sigma_s^2(\nu)$; a small spike could make the crossover explicit (OU-process shared bias, exactly solvable).
 - **Loewner-order statement.** The tensor versions of the deviation and the common-source bound are asserted from the same arguments but not written out; mechanical, worth landing when a tensor-tempo consumer needs them.
 - **Attribution conventions.** If a downstream segment ever needs per-channel corrected contributions (e.g. for channel-investment decisions in TST), Shapley/sequential attribution over the ordered filter is the natural declared-convention candidate — flagged so the convention-dependence is chosen consciously, not rediscovered.
 - **Provenance.** Derived in spikes/spike-adaptive-tempo-redundancy-penalty-2026-07-15.md (full trail, counterexample arithmetic, ripple-site record). Ripple audit executed 2026-07-16 (full-ingestion pass incl. segment git histories): disc-independence-audit §3 repaired (2026-07-15, with the landing); der-team-persistence corrected (false upper-bound + MI-penalty claims deleted, two-sided scope + triangulation stated, depends added); def-strategic-tempo scope made explicit in Epistemic Status (no false claim existed there); der-tempo-composition WN corrected two-sidedly — sharper finding: the synergistic direction conditions the sub-additivity bound $\mathcal{T}_c \leq \sum \mathcal{T}_i$ itself (fusion over pooled raw observations can exploit cross-agent noise correlation), flagged there for a follow-on derivation; hyp-communication-gain corroboration WN re-pointed; the CMI-subtracting figure candidate in def-adaptive-tempo WN annotated as N1-refuted-as-drawn.
-

Appendix AT

Derivation: Matrix-Loewner Persistence Condition

Under matrix adaptive tempo $\mathcal{T}$ from #def-adaptive-tempo's Tensor extension (per-channel primitive $K^{(k)} = (H_M + H_L^{(k)})^{-1}H_L^{(k)}$ from #deriv-fisher-local-update-gain) and matrix disturbance covariance Σ_w , the operational persistence condition lifts from scalar to matrix-Loewner form: the agent persists iff $\mathcal{T}$ is Hurwitz (structural persistence) and the stationary covariance Σ_∞ — solving the continuous Lyapunov equation $\mathcal{T}\Sigma_\infty + \Sigma_\infty\mathcal{T}^T = \Sigma_w$ — is strictly Loewner-dominated by the diagonal of squared per-direction critical thresholds $D_\delta := \text{diag}(\delta_{\text{critical},k}^2)$. The matrix-Loewner condition $\Sigma_\infty < D_\delta$ recovers #result-persistence-condition's scalar Model S form as the isotropic special case and #result-per-dimension-persistence's per-coordinate form as the diagonal special case. Crucially, when $\mathcal{T}$ has off-diagonal entries that misalign with the coordinate basis — the generic situation under cross-dimensional correction — the per-coordinate condition is *unsafe*: it can declare persistence on systems whose stationary mismatch will exceed task-adequacy along the diagonal direction. The matrix-Loewner form is the canonical anisotropic persistence condition; per-coordinate is its diagonal- $\mathcal{T}$ -axis-aligned- D_δ special case, sharp where the coordinate basis happens to be the correction-machinery's eigenbasis and unsafe outside it.

The framework's posture is sharp: **scalar and per-coordinate forms are special cases; the matrix form is the load-bearing object the others project from.** A constructive 2D counterexample (symmetric $\mathcal{T}$ with cross-coordinate correction, isotropic disturbance and threshold) makes the unsafety concrete — per-coordinate evaluates the wrong directions, returning PASS where matrix-Loewner correctly returns FAIL with the failing direction running 45° off-axis. The result composes with the per-direction information-rate floor: no Shannon-type aggregation across dimensions compensates for one direction being capacity-starved (*bandwidth-per-direction is the binding quantity, not total bandwidth*). The cleanest current application is multi-modal AI pipelines with asymmetric per-modality channel budgets, where aggregate-rate metrics systematically miss directional failures the matrix-Loewner reading catches.

Formal Expression.

Setup — Model S with matrix correction.

Consider linear stochastic mismatch dynamics

Definition (matrix-Model-S-mismatch-dynamics).

$$d\delta_t = -\mathcal{T} \delta_t dt + \Sigma_w^{1/2} dW_t \tag{AT.1}$$

with $\delta_t \in \mathbb{R}^d$, matrix adaptive tempo $\mathcal{T} \in \mathbb{R}^{d \times d}$, disturbance covariance $\Sigma_w > 0$, and standard d -dimensional Brownian W_t . Per #def-adaptive-tempo's Tensor extension, $\mathcal{T} = \sum_k \nu^{(k)} K^{(k)}$ aggregates the per-channel matrix gain operators from #deriv-fisher-local-update-gain.

Stationary covariance — continuous Lyapunov equation.

The stationary distribution of the linear stochastic dynamics is Gaussian $\mathcal{N}(0, \Sigma_\infty)$, with Σ_∞ the unique positive-definite solution of the **continuous Lyapunov equation**

Derived (stationary-covariance, exact under Hurwitz $\mathcal{T}$ and $\Sigma_w > 0$).

$$\boxed{\mathcal{T} \Sigma_\infty + \Sigma_\infty \mathcal{T}^T = \Sigma_w.} \quad (\text{AT.2})$$

Existence and uniqueness of $\Sigma_\infty > 0$ are equivalent to $\mathcal{T}$ being **Hurwitz** — all eigenvalues having strictly positive real part. The closed-form integral representation is

$$\Sigma_\infty = \int_0^\infty e^{-\mathcal{T}t} \Sigma_w e^{-\mathcal{T}^T t} dt. \quad (\text{AT.3})$$

When $\mathcal{T}$ is symmetric and commutes with Σ_w , the solution simplifies to $\Sigma_\infty = \frac{1}{2} \mathcal{T}^{-1} \Sigma_w$. Otherwise the integral form (or any standard Lyapunov-equation solver) gives the stationary covariance.

The matrix-Loewner persistence condition.

The agent persists operationally iff:

Derived (matrix-persistence-loewner-Model-S).

(MP-1) Structural persistence. $\mathcal{T}$ is Hurwitz — equivalently, Σ_∞ exists as the unique positive-definite solution of the Lyapunov equation. This is the matrix analog of #result-persistence-condition's structural condition $\alpha > 0$ (Hurwitz reduces to $\mathcal{T}_0 > 0$ in the scalar case).

(MP-2) Matrix-Loewner task adequacy.

$$\boxed{\Sigma_\infty < D_\delta := \text{diag}(\delta_{\text{critical},k}^2)} \quad (\text{AT.4})$$

with $<$ the strict Loewner (positive-definite) order: $D_\delta - \Sigma_\infty > 0$. The stationary mismatch ellipsoid is contained strictly inside the per-direction critical-threshold ellipsoid.

The two conditions are independent in the same sense as #result-persistence-condition's structural-vs-task-adequacy split: (MP-1) is a property of the correction machinery alone; (MP-2) adds the domain-specific bound on per-direction tolerable mismatch.

Three equivalent restatements. The Loewner condition (MP-2) admits three equivalent forms:

(MP-2a) Per-direction. For every unit direction $\hat{v} \in \mathbb{R}^d$:

$$\hat{v}^T \Sigma_\infty \hat{v} < \hat{v}^T D_\delta \hat{v} = \sum_k v_k^2 \delta_{\text{critical},k}^2. \quad (\text{AT.5})$$

The projected stationary variance along $\hat{v}$ must be less than the direction-projected squared threshold along the same direction.

(MP-2b) Spectral. The generalized eigenvalue problem $\Sigma_\infty x = \lambda D_\delta x$ satisfies $\lambda_{\max} < 1$, equivalently

$$\lambda_{\max}(D_\delta^{-1/2} \Sigma_\infty D_\delta^{-1/2}) < 1. \quad (\text{AT.6})$$

The principal eigenvector identifies the *worst direction*; the eigenvalue's distance from 1 is the worst-direction safety margin.

(MP-2c) Ellipsoid containment. The stationary noise ellipsoid $\{x : x^T \Sigma_\infty^{-1} x \leq 1\}$ is contained strictly inside the threshold ellipsoid $\{x : x^T D_\delta^{-1} x \leq 1\}$.

The three are mathematically equivalent under standard matrix-analysis (Horn & Johnson 2013, §4.6 and §7). (MP-2a) is cleanest for hand reasoning, (MP-2b) for computation (single generalized-eigenvalue solve), (MP-2c) for geometric intuition.

Recovery of existing AAT persistence forms.

Table AT.1

<i>Special case</i>	<i>(MP-2) reduces to</i>	<i>Matches</i>
Isotropic ($\mathcal{T} = \mathcal{T}_0 I$, $\Sigma_w = \sigma_w^2 I$, $\delta_{\text{critical}} = \delta_0 1$)	$\sigma_w^2 / (2\mathcal{T}_0) < \delta_0^2$, i.e., $\mathcal{T}_0 > \sigma_w^2 / (2\delta_0^2)$	Scalar Model S form, #result-persistence-condition
Diagonal $\mathcal{T}$, Σ_w in the coordinate basis of D_δ	$\sigma_{w,k}^2 / (2\mathcal{T}_{kk}) < \delta_{\text{critical},k}^2$ per coordinate k	Per-coordinate Model S RMS, #result-per-dimension-persistence
Symmetric $\mathcal{T}$ commuting with Σ_w , general D_δ	Per-eigendirection inequality in $\mathcal{T}$'s eigenbasis, with direction-projected thresholds	New content recovered from the matrix form — extends per-coordinate to non-axis-aligned $\mathcal{T}$

The third row is the content #result-per-dimension-persistence does not currently carry: when $\mathcal{T}$ and Σ_w share an eigenbasis that is *not* the coordinate basis of D_δ , the per-coordinate form misses the right direction-projected thresholds and the matrix-Loewner form handles them correctly.

Where per-coordinate is unsafe — a constructive counterexample. The deeper structural question: is the matrix-Loewner form merely a generalization that agrees with per-coordinate in all real cases, or is it *strictly sharper* — does there exist a regime where per-coordinate says PASS but matrix-Loewner correctly says FAIL? The construction below establishes the latter: the per-coordinate form is **unsafe** under cross-dimensional correction.

Construction. Take $d = 2$, $\Sigma_w = I$, and

$$\mathcal{T} = \begin{pmatrix} 1 & -0.9 \\ -0.9 & 1 \end{pmatrix}. \quad (\text{AT.7})$$

This $\mathcal{T}$ is symmetric positive-definite with eigenvalues 1.9 (along $(1, -1)/\sqrt{2}$) and 0.1 (along $(1, 1)/\sqrt{2}$) — strongly anisotropic correction, with the $(1, 1)$ direction the weak axis. Per-coordinate naive reading of “diagonal of $\mathcal{T}$ ” sees uniform $\mathcal{T}_{11} = \mathcal{T}_{22} = 1$ and concludes correction is uniform at rate 1 on each coordinate.

Stationary covariance. Since $\mathcal{T}$ is symmetric, $\Sigma_\infty = \frac{1}{2}\mathcal{T}^{-1}\Sigma_w = \frac{1}{2}\mathcal{T}^{-1}$. Direct computation:

$$\Sigma_\infty = \frac{1}{0.38} \begin{pmatrix} 1 & 0.9 \\ 0.9 & 1 \end{pmatrix} \approx \begin{pmatrix} 2.63 & 2.37 \\ 2.37 & 2.63 \end{pmatrix}, \quad (\text{AT.8})$$

with eigenvalues 5.00 along $(1, 1)/\sqrt{2}$ and 0.26 along $(1, -1)/\sqrt{2}$.

Set $\delta_{\text{critical}} = (1.7, 1.7)$, so $D_\delta = 2.89 I$.

- *Per-coordinate check* (#result-per-dimension-persistence *form*): per-coordinate stationary variance is $\Sigma_{\infty, kk} = 2.63 < 2.89 = \delta_{\text{critical}, k}^2$ for each k . **Per-coordinate says PASS.** ✓
- *Matrix-Loewner check* (MP-2 *here*): worst-direction projected variance is $\lambda_{\max}(\Sigma_\infty) = 5.00$ along $(1, 1)/\sqrt{2}$. Direction-projected squared threshold along the same direction is $\hat{v}^T D_\delta \hat{v} = 2.89$. The condition $5.00 < 2.89$ fails. Equivalently, $\lambda_{\max}(D_\delta^{-1/2} \Sigma_\infty D_\delta^{-1/2}) = 5.00/2.89 \approx 1.73 > 1$. **Matrix-Loewner says FAIL.** ✗

Conclusion. Per-coordinate gives a false-pass. The agent will routinely produce mismatch vectors with RMS magnitude $\sqrt{5.00} = 2.24$ along the $(1, 1)/\sqrt{2}$ direction, exceeding the direction-projected threshold of $\sqrt{2.89} = 1.70$ on that direction. The per-coordinate analysis missed the diagonal direction because each coordinate alone stays within $\delta_{\text{critical}, k} = 1.7$ — the failure is *only* visible at the matrix-Loewner level. Cross-dimensional correction (the off-diagonal entries of $\mathcal{T}$) is what creates the diagonal-direction concentration that per-coordinate cannot detect.

Epistemic Status.

Conditional. Max attainable: *exact* under (i) Model S linear stochastic dynamics, (ii) Hurwitz matrix tempo $\mathcal{T}$, (iii) positive-definite disturbance covariance Σ_w , (iv) diagonal critical-threshold structure $D_\delta = \text{diag}(\delta_{\text{critical}, k}^2)$. The Lyapunov equation, the existence-iff-Hurwitz fact, and the matrix-Loewner equivalence forms are all standard from Bellman 1960 / Khalil 2002 / Horn-Johnson 2013; the AAT-distinctive content is the application to the persistence condition and the recognition that per-coordinate is strictly sharper-or-unsafe rather than uniformly sharper.

What is load-bearing:

- The matrix-Loewner form as the canonical anisotropic persistence condition, with scalar and per-coordinate forms as special cases recovered exactly.
- The strict-sharpness result via the §4 counterexample: per-coordinate gives a false-pass when $\mathcal{T}$ has off-diagonal entries that misalign with the coordinate basis. The per-coordinate form is not merely incomplete in this regime — it is unsafe.
- The three equivalent restatements (per-direction, spectral, ellipsoid-containment) with their different operational readings.
- The structural-vs-task-adequacy decomposition from #result-persistence-condition transfers directly: (MP-1) is the matrix structural condition (Hurwitz); (MP-2) is the matrix task adequacy (Loewner $\Sigma_\infty < D_\delta$). The two conditions remain independent — Hurwitz $\mathcal{T}$ does not imply $\Sigma_\infty < D_\delta$ for any particular D_δ , and conversely.

What is not established here:

- The nonlinear matrix-sector extension. Under sector-bounded $F(\mathcal{T}, \delta)$ with $\delta^T F \geq \alpha I$ in matrix-Loewner form, the matrix-Lyapunov machinery extends to give an ultimate-bound matrix; the linear case here is the template instantiation. Full matrix-sector treatment is follow-on work in #deriv-sector-condition.
- The Model D matrix lift. Under deterministic bounded disturbance $\|w_t\|_M \leq \rho$ in a quadratic norm, the ultimate-bound becomes an ellipsoid containment problem (a linear matrix inequality). Mechanical lift; not derived here.
- The non-Hurwitz boundary behaviour (eigenvalue exactly on the imaginary axis). Outside Hurwitz, the linear analysis fails — no stationary distribution. This matches the scalar case's $\alpha > 0$ requirement.
- The promotion of #result-adversarial-tempo-advantage to matrix form. The matrix-Loewner form sharpens the adversarial result via worst-direction targeting (an adversary controlling Σ_w^{adv} maximizes $\lambda_{\max}(D_\delta^{-1/2} \Sigma_w^{\text{adv}} \mathcal{T}_{\text{eff}}^{-1} D_\delta^{-1/2})$), but the full adversarial promotion is a separate cycle.
- The matrix extension of #deriv-persistence-cost's information-rate floor. The natural lift is $\dot{R} \geq \frac{1}{2} \text{Tr}(\mathcal{T})$ — per-eigendirection application of the scalar derivation. Worth a cross-reference Working Note in that segment if it becomes load-bearing; not derived here.

Discussion.

Why the Loewner form is the natural matrix lift. The Lyapunov equation $\mathcal{T}\Sigma_\infty + \Sigma_\infty\mathcal{T}^T = -\Sigma_w$ is itself a matrix-Loewner statement: it says Σ_∞ balances correction (via $\mathcal{T}$) against disturbance (via Σ_w) in the PSD ordering. The task-adequacy condition $\Sigma_\infty < D_\delta$ reads this balance against the threshold ellipsoid. The whole structure lives in the cone of positive-semidefinite matrices ordered by Loewner dominance; scalar persistence is its 1D shadow, per-coordinate persistence is its axis-aligned-diagonal shadow, and the matrix form is the load-bearing object the shadows project from.

The structural-vs-task-adequacy decomposition transfers. #result-persistence-condition's split — structural persistence (machinery works) vs task adequacy (machinery works well enough) — survives intact in the matrix lift: Hurwitz $\mathcal{T}$ is the matrix structural condition; Loewner $\Sigma_\infty < D_\delta$ is the matrix task-adequacy condition. The two are independent; an agent can have Hurwitz $\mathcal{T}$ (correction works in every direction) but fail Loewner on the worst direction (correction is not fast enough on the weak axis for the per-direction threshold). The remedies are different: structural

failure requires changing the correction architecture (#result-structural-adaptation-necessity); task-adequacy failure requires either increasing $\mathcal{T}$ on the weak direction, decreasing Σ_w on the weak direction, or relaxing the threshold along that direction.

The weak-direction-bottleneck argument transfers and sharpens. #result-per-dimension-persistence’s key insight — the weak dimension is the bottleneck — generalizes to the weak *direction* in the matrix lift: the worst eigenvector of Σ_∞ relative to D_δ is what gates persistence. When this eigenvector aligns with a coordinate axis, the per-coordinate analysis catches it. When it points off-axis — the generic situation under cross-dimensional correction — the per-coordinate analysis misses it. The §4 counterexample is the canonical illustration: the worst direction is $(1, 1)/\sqrt{2}$, off-axis from both coordinate directions; per-coordinate evaluates the wrong directions; matrix-Loewner evaluates the right ones.

Cross-coordinate correction is the generic case under Fisher-local invariance. The matrix gain operator $K = (H_M + H_L)^{-1}H_L$ from #deriv-fisher-local-update-gain is diagonal in the coordinate basis only when prior precision H_M and observation Fisher H_L share an eigenbasis aligned with the coordinate axes. In general — and certainly in the natural-gradient regime where the agent’s parameterization is statistical-manifold-natural rather than user-chosen — they do not, and K has off-diagonal entries. The matrix-Loewner condition becomes the canonical form for any Fisher-local-derived persistence analysis; the per-coordinate form is the exceptional case where prior, likelihood, and threshold all happen to share the coordinate axes.

Cross-reference to existing AAT machinery. The matrix-Loewner form composes with:

- #result-sector-persistence-template (T1)–(T3) as the linear-case template instantiation; the matrix sector inequality $\delta^T F + F^T \delta \geq 2\alpha I$ would extend the matrix-Lyapunov result to nonlinear F via the template’s (T2) condition lifted to matrix form.
- #deriv-sector-condition Prop A.1S (the Lyapunov-stochastic derivation underlying #result-persistence-condition’s scalar form) — the matrix derivation here is the analog at the matrix-Lyapunov level.
- #deriv-fisher-local-update-gain (matrix gain K as the per-channel primitive) and #def-adaptive-tempo (Tensor extension aggregating $\mathcal{T} = \sum_k \nu^{(k)} K^{(k)}$).

On the scope tag downstream. Until the broader promotion of #result-persistence-condition, #result-adversarial-tempo-advantage, #form-composition-closure, #der-team-persistence, #deriv-critical-mass-composition to invoke the matrix form directly, those segments continue to read scalar and inherit the “scalar / isotropic / nonredundant-channel scope” tag from #def-adaptive-tempo’s Epistemic Status. The per-direction primitive is in place; the downstream promotion is a follow-on cycle in TODO.md.

Working Notes.

- **Counterexample as load-bearing content.** Section 4’s $\mathcal{T} = \begin{pmatrix} 1 & -0.9 \\ -0.9 & 1 \end{pmatrix}$ example is the segment’s strongest content — it shifts the matrix-Loewner form from “interesting generalization” to “the safe condition; per-coordinate is unsafe under cross-dimensional correction.” The example uses the smallest nontrivial dimension ($d = 2$), the simplest non-diagonal symmetric $\mathcal{T}$ structure, and isotropic disturbance / threshold parameters. Generalizations to $d > 2$, asymmetric $\mathcal{T}$, and anisotropic Σ_w/D_δ are mechanical; the qualitative phenomenon (off-diagonal correction $\rightarrow$ diagonal-direction variance concentration $\rightarrow$ per-coordinate false-pass) is robust.
- **Empirical validation.** A 2D simulation matching the §4 parameters would confirm the predicted false-pass behaviour of per-coordinate: run the linear SDE with $\mathcal{T}, \Sigma_w$ from §4 for N steps; observe that the per-coordinate marginals stay within ± 1.7 most of the time (per-coordinate “passes”) while the diagonal-direction projection routinely exceeds ± 1.7 (matrix-Loewner “fails”). Mechanical simulation; would land cleanly under obs-section-i-validation-simulations as a new variant.

- **Open work — nonlinear matrix-sector.** The matrix-Loewner form is derived for linear correction. The nonlinear matrix-sector lift would replace the Lyapunov equation with the matrix-sector inequality $F(\delta) + F(\delta)^T \succeq 2\alpha\delta\delta^T/\|\delta\|^2$ (or a matrix-sector form derived from a quadratic Lyapunov function). Composes with #deriv-sector-condition's nonlinear scalar machinery via the matrix template; not derived here.
 - **Open work — adversarial extension.** The adversarial advantage exponent in #result-adversarial-tempo-advantage lifts to a matrix-eigenvalue-ratio exponent: an adversary controlling Σ_w^{adv} maximizes the bad-direction stationary variance via the matrix-Loewner form. Worth a follow-on derivation if #result-adversarial-tempo-advantage is promoted to matrix form.
 - **Open work — composition lift.** Composite stationary covariance Σ_∞^c solves the Lyapunov equation with composite $\mathcal{T}^c$ and Σ_w^c (aggregated from sub-agents); composite Loewner condition $\Sigma_\infty^c \prec D_\delta^c$ governs composite persistence. Promoting #form-composition-closure, #der-team-persistence, and #deriv-critical-mass-composition to invoke the matrix form is the natural next AAT-1 follow-on cycle.
 - **Landing context.** Landed in the 2026-05-12 AAT-1 follow-on cycle (matrix-Loewner persistence, succeed-beyond-claim); see CHANGELOG 2026-05-12 (late). The load-bearing derivation and counterexample are above; the follow-on extensions (Model D matrix lift, adaptive-gain matrix dynamics, variational matrix form) are flagged in the “Open work” notes here and in TODO. Originating spike is absorbed archaeology, not a live reference.
-

Appendix AU

Derivation: The Survival-Imperative Exploration Drive

The unified policy objective in `#disc-ci-y-unified-objective` combines pragmatic value Q_O with an epistemic exploration term (Causal Information Yield, CIY) weighted by $\lambda_{\text{info}} \propto U_M$. The logic is intuitive: explore when uncertain.

However, AAT’s underlying Lyapunov dynamics (`#result-persistence-condition`) force a *second, structurally distinct* exploration drive: the **Survival Imperative**. This segment derives that an agent must actively seek low-noise (high CIY) observations not just to learn, but to physically survive environmental drift when its model is *confident*.

AAT therefore carries **two parallel exploration drives at opposite ends of the uncertainty spectrum**: the standard epistemic drive scaling as $\lambda_{\text{info}} \propto U_M$ (explore when uncertain — dominant when model uncertainty is high), and the survival drive scaling as $\lambda_{\text{surv}} \propto 1/U_M$ (explore when confident in a drifting world — dominant when model uncertainty is low). The result *bypasses the dark-room critique of active inference* structurally rather than parametrically: an AAT agent in a drifting world physically cannot stop exploring because exploration is forced by the same Lyapunov machinery that grounds the rest of Part II, not by a tunable preference parameter. The operational prescription — *confident agents are not safe agents in drifting environments* — is the formal grounding for the “productive paranoia” heuristic about not resting on past success. The scalar form here admits a “blank-wall” attack resolved in AV by lifting to a Linear Matrix Inequality on the Fisher Information Matrix with a positive-semidefinite matrix Lagrange multiplier; the dual-drive picture survives the lift in matrix form.

The Persistence-Bounded Objective.

The agent’s primary goal is to maximize pragmatic value: $\max_{\pi} \mathbb{E}_{a \sim \pi}[Q_O(a)]$

The agent is subject to the structural persistence condition ($\alpha R > \rho^{\text{eff}}$), meaning it only survives if its steady-state mismatch $R^* < R$. Under Model D dynamics, $R^* = \rho^{\text{eff}}/(\eta^* c_{\min})$. Using the optimal update gain $\eta^* = \frac{U_M}{U_M + U_o}$, we can rewrite the survival constraint exactly: $\frac{U_M}{U_M + U_o} > \frac{\rho^{\text{eff}}}{R c_{\min}}$

Solving for U_o , we find the maximum observation noise the agent can tolerate without destabilizing: $U_o^{\max} = U_M \left(\frac{R c_{\min}}{\rho^{\text{eff}}} - 1 \right)$

This creates a strict constraint on the agent’s action selection: $\mathbb{E}_{a \sim \pi}[U_o(a)] \leq U_o^{\max}(M_t)$

The Tragedy of the Confident Agent.

Notice the proportionality: $U_o^{\max} \propto U_M$.

As $U_M \rightarrow 0$ (the agent becomes highly confident), $U_o^{\max} \rightarrow 0$. The survival constraint becomes infinitely tight.

If the environment is drifting ($\rho^{\text{eff}} > 0$), a confident agent inherently ignores noisy observations ($\eta^* \rightarrow 0$). To survive the drift, the agent *must* force an update. The only way to penetrate its own stubbornness is to generate an observation so pristine (low U_o) that the agent is forced to correct its course. Thus, a confident agent in a drifting world is *forced* to explore (seek low U_o) simply to survive.

Conversely, an uncertain agent (high U_M) has a loose constraint (large $U_o^{\max}$). It updates easily ($\eta^* \rightarrow 1$), so it can safely exploit noisy states without fear of structural collapse.

Deriving the Survival Lagrange Multiplier λ_{surv} .

Formulating the KKT Lagrangian for this constrained optimization: $\mathcal{L} = \mathbb{E}[Q_O(a)] - \lambda' \cdot (\mathbb{E}[U_o(a)] - U_o^{\max})$

Derivation (survival-exploration).

The shadow price λ' scales as $1/(U_o^{\max} - \mathbb{E}[U_o])$. At the boundary, $\lambda' \propto 1/U_o^{\max}$, yielding: $\lambda' \propto \frac{\rho^{\text{eff}}}{(Rc_{\min} - \rho^{\text{eff}}) \cdot U_M}$

To align this with the CIY objective, we invoke the structural assumption that Causal Information Yield is inversely proportional to observation ambiguity: $\text{CIY}(a) \propto 1/U_o(a)$. This scalar reciprocal mapping is a heuristic placeholder: the mathematically clean correspondence between the survival-imperative Lagrangian and the CIY/EIG objective is the LMI / matrix-valued lift in `#deriv-causal-ib-lmi`, where the Fisher Information Matrix and a positive-semidefinite Lagrange multiplier Λ replace the scalar substitution below. Readers tracking exactness rather than directional fidelity should treat the rest of this section as the scalar shadow of that LMI derivation; the *honest limit* sub-section below and `#deriv-causal-ib-lmi` together carry the formal load.

Transforming the Lagrangian penalty from U_o to $1/U_o$ (which introduces a non-linear mapping of the multiplier, operating-point dependent), the equivalent survival weight on CIY becomes: $\lambda_{\text{surv}}(M_t) \propto \frac{\rho^{\text{eff}}}{(Rc_{\min} - \rho^{\text{eff}}) \cdot U_M}$

Epistemic Status.

Derived (conditional). The derivation of the $U_o^{\max}$ bound and the KKT multiplier λ' is exact, conditional on the continuous-time Lyapunov bounds from `#result-sector-condition-stability`. The mapping to $\lambda_{\text{surv}} \cdot \text{CIY}$ is *heuristic*, conditional on the assumed structural relationship $\text{CIY}(a) \propto 1/U_o(a)$ and the non-linear transformation of the shadow price.

Discussion.

Two Parallel Exploration Drives. AAT dictates two correlated but distinct motivations for exploration, acting at opposite ends of the uncertainty spectrum: 1. **Epistemic Information Gain** ($\lambda_{\text{info}} \propto U_M$): Explored in `#disc-ciy-unified-objective`. The agent explores to reduce its model uncertainty. Dominates when U_M is high. 2. **Lyapunov Survival Imperative** ($\lambda_{\text{surv}} \propto 1/U_M$): Derived here. The agent explores to force high-fidelity corrections against environmental drift. Dominates when U_M is low.

Resolving the IB β vs ρ conflation. This derivation clarifies a structural tension in the standard IB formulation (`#form-information-bottleneck`). Standard IB natively discards old information as ρ rises because the joint distribution $P(\mathcal{C}_t, O_{t+1})$ decorrelates. The agent's *representational preference* (β) does not need to change. However, the agent's *action policy* is tightly coupled to ρ . As ρ increases, the survival constraint tightens ($U_o^{\max} \downarrow$), forcing the Lagrange multiplier λ_{surv} to spike. The environment dictates the representation natively; survival dictates the policy.

Connection to Active Inference. This derivation completely sidesteps the Active Inference “preferences-as-priors” assumption (the dark room problem). The survival exploration drive exists purely because failing to bound U_o mathematically guarantees structural collapse ($\alpha R < \rho^{\text{eff}}$). Exploration is derived as a physical survival imperative, not a psychological preference.

Honest limit: the scalar formulation admits a “blank wall” attack. The survival constraint as derived bounds the *scalar* observation-noise expectation $\mathbb{E}_\pi[U_o(a)] \leq U_o^{\text{max}}$ — it cares only about the magnitude of U_o , not its alignment with the directions in which ρ^{eff} is acting. In multidimensional environments where some coordinates of the drift are inactive, an action that selects an observation channel orthogonal to the drift (a constant signal — the “blank wall”) drives U_o to near-zero in that subspace and satisfies the scalar survival condition, while the agent’s prediction error in the drifting subspace grows without bound because the observations carry no information about the drifting coordinates. The scalar reduction collapses the dimensionality that matters.

The scalar derivation here remains correct in its scope (1D environments and isotropic-disturbance multidimensional environments where any low-noise observation suffices), and the qualitative dual-drive picture survives the lift to higher dimensions. The full multidimensional resolution — upgrading the scalar inequality to a Linear Matrix Inequality on the Fisher Information Matrix, with a positive-semidefinite matrix Lagrange multiplier Λ that distinguishes by direction rather than magnitude — is in `#deriv-causal-ib-lmi`.

Findings.

Survival-Imperative Exploration as Lyapunov-Forced Drive. Brief: A confident agent in a drifting world is mathematically forced to seek high-fidelity (low-noise) observations not because exploration is informationally valuable but because failing to do so violates the structural persistence condition and guarantees collapse. The maximum tolerable observation noise U_o^{max} scales with U_M — meaning the more confident the agent becomes, the tighter the survival-noise constraint becomes, and the more the agent is forced to actively seek low-ambiguity observations to penetrate its own stubbornness. The result identifies a second exploration drive ($\lambda_{\text{surv}} \propto 1/U_M$) structurally distinct from and complementary to the standard epistemic information-gain drive ($\lambda_{\text{info}} \propto U_M$), with the two acting at opposite ends of the uncertainty spectrum.

Impact: Provides AAT with a derivation of *why exploration is required* that does not route through epistemic-value-of-information arguments and therefore does not inherit the dark-room-problem objection that has dogged active-inference accounts. The survival imperative falls out of the same Lyapunov persistence condition that grounds the rest of Part II — exploration is forced by the same mathematics that forces the agent’s adaptive responsiveness, not added as a separate motivational layer. The dual-drive picture also resolves a structural tension in standard Information Bottleneck formulations between the representational β and the environmental drift rate ρ : the environment dictates representation natively (IB’s joint distribution decorrelates as ρ rises); survival dictates policy (the constraint tightens as ρ rises). The scalar form admits a “blank wall” attack, resolved in `#deriv-causal-ib-lmi` by lifting the constraint to an LMI on the Fisher Information Matrix; the dual-drive picture survives the lift in matrix form.

Novelty Claim: *Claim differentiation* on the structural source of agentic exploration. Active-inference work derives exploration from epistemic value inside expected free energy; bounded-rationality and information-cost frameworks derive exploration from utility-shaped information processing. This finding derives exploration from Lyapunov-persistence collapse — a structural-stability argument rather than a value-of-information argument — and identifies a second exploration multiplier inversely proportional to model uncertainty, complementary to the epistemic drive that scales directly with it.

Related Work:

- Friston, Rigoli, Ognibene, Mathys, FitzGerald & Pezzulo 2015, “Active inference and epistemic value,” *Cognitive Neuroscience* 6:187–214 (published 2015, found 2026-04 via Undermind Pillar 1 report) — *adjacent literature with structural contrast* — derives exploration from epistemic value inside expected free energy; the present finding’s exploration drive is forced by Lyapunov bounds rather than preferences-as-priors, sidestepping the dark-room objection.
- Friston, Samothrakis & Montague 2012, “Active inference and agency: optimal control without cost functions,” *Biological Cybernetics* 106:523–541 (published 2012, found 2026-04) — *adjacent literature* — broader active-inference framing within which the EFE epistemic-value account sits; same structural-vs-preference contrast.
- Ortega & Braun 2012, “Thermodynamics as a theory of decision-making with information-processing costs,” doi:10.1098/rspa.2012.0683 (published 2012, found 2026-04) — *adjacent literature* — bounded-rationality framework deriving exploration-like behavior from information-processing costs; the present finding derives a structurally distinct second drive (Lyapunov survival) without inheriting the bounded-rationality objective form.
- Genewein, Leibfried, Grau-Moya & Braun 2015, “Bounded Rationality, Abstraction, and Hierarchical Decision-Making,” *Frontiers Robotics AI* 2:27 (published 2015, found 2026-04) — *adjacent literature* — hierarchical extension of Ort12c; same structurally-different motivational source.
- Tishby, Pereira & Bialek 2000, “The information bottleneck method” (published 2000, found 2024) — *formal antecedent* — the standard IB formulation against which the β -vs- ρ tension is resolved; the survival imperative provides the missing policy-side coupling that standard IB does not natively carry.

Search Log:

- 2026-04 (*nominally comprehensive on the active-inference / bounded-rationality contrast*, via ref/Novelty_defense_and_integration.md Pillar 1): The Undermind report identified Friston 2015 as the canonical EFE-epistemic-value account against which this finding’s structural-stability source for exploration differentiates, and the Ortega/Genewein bounded-rationality line as the closest information-cost-shapes-policy precursor. The verdict on Pillar 1 as a whole was *Conceptual Precursor* (High confidence) for the no-go half (#der-causal-insufficiency-detection); the survival-imperative half is the constructive complement, framed as differentiation from epistemic-value exploration rather than as a no-go contribution.
- 2025 (*intuition-only*): Pre-Undermind expectation was that the closest priors would be in active-inference (Friston) and Information Bottleneck (Tishby) literatures; this was confirmed and refined by the comprehensive search.

Working Notes.

- **Empirical validation.** The derivation is validated in `empirica/track-b-nonlinear/variants/variant_causal_ib.py` (results: `variant_causal_ib_results.md`). In a highly volatile environment ($\rho = 0.5$) where exploitation produces ambiguous observations ($U_o = 100.0$), a purely greedy agent ($\lambda = 0$) suffers a 0% survival rate because its update gain collapses.
- **LMI generalization landed.** The multidimensional resolution is in `#deriv-causal-ib-lmi` (matrix Lagrangian on the Fisher Information Matrix; blank-wall attack resolved by complementary slackness on direction). Spike trail: `spikes/.integrated/spike-causal-information-bottleneck.md` §7 (critique) and `spikes/.integrated/spike-causal-ib-lmi-repair.md` (LMI derivation).
- **Cross-reference to NeurIPS Paper 1.** The survival-imperative drive is formalized at theorem grade in NeurIPS 2026 Paper 1 (“Tragedy of the Confident Agent: Forced Exploration for Survival in Drifting Environments”, `~/src/neurips/01-tragedy-confident-agent/`): (i) §3 Lemma 2.1(ii) supplies the worst-case finite **exit-time bound** $T_{\text{exit}}(\delta_0) \leq \frac{1}{\alpha+\gamma} \log \frac{\rho/(\alpha+\gamma)-\|\delta_0\|}{\rho/(\alpha+\gamma)-R}$

that the controller-design overshoot is justified against; (ii) §5 Mechanism Step 4 + Appendix-KKT resolves the **shadow-price / controller-divergence puzzle** — the static-program LP multiplier is finite, while the survival-margin controller's λ_{surv} divergence is a controller-design overshoot ensuring loop closure inside the exit-time window (divergence is on the environment-side clock, not the controller-side multiplier); (iii) §App-A.7 / #prop-fagp names the **F-A-G-P enforcement framework** (strict-Feasibility / drift-Aligned bonus / Gain dominance / argmax-Presentable) as the four-condition operational decomposition for valid survival-margin controller family members. See spikes/neurips-back-integration-2026-05-08.md §1 Paper 1 entries 1, 3, 4 for the cross-mapping.

Appendix AV

Derivation: Linear Matrix Inequality Form of the Causal-IB Survival Bound

In `#deriv-causal-ib-exploration`, the survival-imperative exploration drive was derived in scalar form: $\text{bound } \mathbb{E}_\pi[U_o(a)] \leq U_o^{\max}$ to keep the agent’s mismatch within survivable bounds. The scalar bound admits a “blank wall” attack — actions that minimize U_o in a subspace orthogonal to the drift (e.g., observing a constant signal uncorrelated with the environmental drift) satisfy the survival math without probing the drifting coordinates, and the agent’s prediction in those subspaces diverges unboundedly. This segment lifts the survival constraint to a Linear Matrix Inequality on the Fisher Information Matrix, replacing the scalar shadow price with a positive-semidefinite matrix Lagrange multiplier Λ that distinguishes by direction, not just by magnitude.

Under the matrix form, the action’s exploration bonus is the trace $\text{Tr}(\Lambda \cdot J_o(a))$, and complementary slackness on the directional constraint *mathematically zeros out the bonus for blank-wall actions*: the bonus structurally requires the agent to probe the drifting directions. The result converts the unified policy objective from heuristic to conditional-theorem-grade under the matrix-LMI form; the structural form is grounded in AAT’s physical survival bounds and standard semidefinite programming rather than treated as an ad-hoc heuristic. The same Fisher Information Matrix appears here in *three* structurally distinct framework loads — credit assignment (AM), coordinate forcing (5.5 4th instance), and now survival — because the framework’s coordinates are forced rather than chosen; this segment is the first explicit cross-instance composition where coordinate-forcing and survival interact within the same machinery. Composition with the persistence-information-rate floor sharpens: the floor names *what bandwidth must be met*; the LMI names *in what directions to spend it*. Without the LMI the bandwidth could be supplied in irrelevant directions; with it, complementary slackness routes the bandwidth toward the drifting modes — a directional version of the central inequality.

Formal Expression.

Multidimensional mismatch dynamics.

The mismatch state $\delta_t \in \mathbb{R}^n$ evolves under linear-Gaussian dynamics with action-dependent observation noise:

Definition (multidim-mismatch).

$$\delta_{t+1} = (I - K_t H) \delta_t + w_t - K_t v_t(a_t) \tag{AV.1}$$

where $w_t \sim \mathcal{N}(0, Q_p)$ is the environmental drift, $v_t \sim \mathcal{N}(0, R_o(a_t))$ is the observation noise for action a_t , and the optimal Kalman gain is $K_t = P_t H^T (H P_t H^T + R_o(a_t))^{-1}$ with P_t the agent’s prior uncertainty covariance.

The steady-state mismatch covariance Σ_δ satisfies the discrete-time algebraic Riccati equation. *Survival* requires the spectral bound

$$\lambda_{\max}(\Sigma_\delta) < R^2 \quad (\text{AV.2})$$

— the matrix lift of the scalar bound $R^* < R$ from #result-persistence-condition.

Fisher Information Matrix as matrix CIY.

For the linear-Gaussian observation model, the action-conditional Fisher Information Matrix is

Definition (matrix-ciy).

$$\mathcal{J}_o(a) = H^T R_o(a)^{-1} H. \quad (\text{AV.3})$$

Where the scalar formulation in #disc-ciy-unified-objective defined CIY as a surrogate satisfying $\text{CIY}(a) \propto 1/U_o(a)$, the multidimensional formulation identifies CIY directly with the action-conditional FIM. The $\text{CIY-}U_o$ surrogate-mapping layer of approximation in the scalar derivation is removed at this layer; the residual CIY-vs-EIG concern (FIM equals EIG-rate only under Gaussian-linear conditions) persists as before.

The information-form Kalman update is additive in the FIM:

$$P_{t+1|t+1}^{-1} = P_{t+1|t}^{-1} + \mathcal{J}_o(a_t), \quad (\text{AV.4})$$

where $P_{t+1|t} = A P_t A^T + Q_w$ is the prediction step that absorbs process noise. The simplification $P_{t+1}^{-1} = P_t^{-1} + \mathcal{J}_o(a_t)$ used in some passes refers to the information-update step alone; the full step adds the prediction contribution. The two coincide once the analysis is in steady state.

LMI survival constraint.

The discrete-time algebraic Riccati equation admits a stabilizing steady-state solution with $\lambda_{\max}(\Sigma_\delta) < R^2$ if and only if the action policy π provides Fisher information that dominates the unstabilizable modes of the system drift. Specifically, the policy must satisfy the Linear Matrix Inequality

Derived (lmi-survival).

$$\mathbb{E}_{a \sim \pi}[\mathcal{J}_o(a)] \succeq \mathcal{J}_{\min}(Q_\rho, A, R^2) \quad (\text{AV.5})$$

where $\mathcal{J}_{\min}$ is the symmetric positive-definite lower bound on the FIM required to keep the steady-state covariance's largest eigenvalue below R^2 . The closed form for $\mathcal{J}_{\min}$ as a function of (Q_ρ, A, R^2) is given by the steady-state DARE condition — standard control-theory machinery (Boyd, Ghaoui, Feron & Balakrishnan 1994 Ch. 3–5; Anderson & Moore 1979). AAT's contribution at this layer is the framing: $\mathcal{J}_{\min}$ enters as the *survival-imperative* Fisher-information floor, not as a controller-design parameter.

Tensor Lagrangian.

The agent maximizes pragmatic value under the LMI constraint:

Derived (tensor-lagrangian).

$$\max_{\pi} \mathbb{E}_{\pi}[Q_O(a)] \quad \text{s.t.} \quad \mathbb{E}_{\pi}[\mathcal{J}_o(a)] \geq \mathcal{J}_{\min}. \quad (\text{AV.6})$$

The Lagrangian for a matrix-inequality-constrained problem uses a positive-semidefinite matrix Lagrange multiplier $\Lambda \succeq 0$:

$$\mathcal{L}(\pi, \Lambda) = \mathbb{E}_{\pi}[Q_O(a)] + \text{Tr}(\Lambda \cdot (\mathbb{E}_{\pi}[\mathcal{J}_o(a)] - \mathcal{J}_{\min})). \quad (\text{AV.7})$$

KKT conditions are primal feasibility ($\mathbb{E}_{\pi}[\mathcal{J}_o(a)] \geq \mathcal{J}_{\min}$), dual feasibility ($\Lambda \succeq 0$), and complementary slackness ($\text{Tr}(\Lambda \cdot (\mathbb{E}_{\pi}[\mathcal{J}_o(a)] - \mathcal{J}_{\min})) = 0$). The optimal action selection rule is

$$a_t \in \arg \max_a [Q_O(a) + \text{Tr}(\Lambda \cdot \mathcal{J}_o(a))]. \quad (\text{AV.8})$$

The trace inner-product $\text{Tr}(\Lambda \cdot \mathcal{J}_o(a))$ is the **matrix exploration bonus** — the multidimensional analog of $\lambda_{\text{surv}} \cdot \text{CIY}(a)$ from the scalar derivation.

Resolution of the blank-wall attack.

By complementary slackness, Λ has support only in eigendirections where the LMI constraint binds — i.e., directions where the agent’s accumulated FIM sits at the threshold $\mathcal{J}_{\min}$. In non-drifting eigendirections of Q_{ρ} , $\mathcal{J}_{\min}$ contributes zero (no information needed there to survive), so Λ has zero weight in those directions.

Derived (blank-wall-resolution).

For a blank-wall action — high $\mathcal{J}_o(a)$ in a non-drifting eigendirection — the trace product $\text{Tr}(\Lambda \cdot \mathcal{J}_o(a))$ evaluates to zero: the action’s FIM eigenvalues lie outside Λ ’s support. The agent receives no exploration bonus for the blank-wall action, even though the action would satisfy the scalar magnitude bound. **The matrix Lagrangian distinguishes by direction, not just by magnitude**, mathematically forbidding the trivial-exploration solutions that the scalar form admitted.

Scalar reduction.

In the 1D case, H and $R_o(a)$ are scalars, $\mathcal{J}_o(a) = R_o(a)^{-1}$, and the LMI degenerates to the scalar inequality $\mathbb{E}_{\pi}[R_o(a)^{-1}] \geq \mathcal{J}_{\min}$ which, with $\mathcal{J}_{\min}^{-1} = U_o^{\max}$, recovers the scalar survival constraint of #deriv-causal-ib-exploration. The matrix multiplier Λ collapses to a scalar shadow price coinciding with λ_{surv} , and the matrix exploration bonus reduces to $\lambda_{\text{surv}} \cdot \text{CIY}(a)$. The multidimensional derivation is structurally consistent with — and properly subsumes — the scalar one.

Derived (scalar-recovery).

Tragedy of the Confident Agent — matrix form.

As the agent's prior covariance P_t shrinks in a particular eigendirection v , the agent's confidence in that direction grows. If v also has support in Q_ρ (a drifting eigendirection), the survival LMI requires $\mathcal{J}_{\min}$ to grow along v to compensate — a confident agent in a drifting direction must source information specifically along that direction, because its low P_t in v means its update gain is small there and slow to absorb new information. By complementary slackness, the matrix shadow price Λ develops a large principal eigenvalue along v , forcing the agent to choose actions whose $\mathcal{J}_o(a)$ has large support on v . The qualitative tragedy survives the lift: confident agents in drifting worlds are mathematically forced to probe the drifting eigendirections specifically, not just to minimize scalar noise.

Derived (matrix-confidence-tragedy).

Epistemic Status.

Derived (conditional). The structural derivation is exact under named conditions:

1. **Riccati stabilizability.** The discrete-time algebraic Riccati equation admits a stabilizing steady-state solution. Standard machinery: Boyd, Ghaoui, Feron & Balakrishnan 1994 §3–5 and Anderson & Moore 1979.
2. **Linear-Gaussian observation model.** The observation $o = Hx + v$ with $v \sim \mathcal{N}(0, R_o(a))$. Non-Gaussian or nonlinear models require local linearization or extended-Kalman treatment.
3. **Information-form Kalman in steady state.** The $\mathcal{J}_o$ -additivity used in the FIM update reflects the information-update step alone; the prediction step's $AP_tA^T + Q_w$ contribution is absorbed into the steady-state Riccati analysis. Transient regimes admit the constraint shape but with $\mathcal{J}_{\min}$ a function of the transient covariance.
4. **DARE-imported $\mathcal{J}_{\min}$.** The closed-form expression for $\mathcal{J}_{\min}(Q_\rho, A, R^2)$ via the DARE is theorem-imported standard control theory. AAT's contribution is the framing — $\mathcal{J}_{\min}$ enters as the survival-imperative FIM floor, with Λ as the directional exploration-bonus shadow price — not the DARE solution itself.

The matrix-CIY mapping ($\mathcal{J}_o(a)$ as multidimensional CIY) is more rigorous than the scalar version: where the scalar formulation needed $\text{CIY} \propto 1/U_o$ as an assumed structural mapping, the matrix formulation directly identifies CIY with the action-conditional FIM, removing one of the two layers of approximation in #disc-ci-y-unified-objective. The CIY-vs-EIG surrogate concern remains: FIM is the EIG-rate for Gaussian linear models but diverges from EIG under non-Gaussian or non-local conditions.

Discussion.

Connection to AAT's existing matrix-form machinery. The LMI repair sits naturally within several existing AAT components rather than being imported from outside:

- #deriv-fisher-whitened-update-rule derives Fisher-information-based credit assignment as AAT-internal under the (PI) parameterization-invariance axiom + Čencov 1982 uniqueness. The FIM appearing in the LMI is the same Fisher metric, used here for survival rather than for credit assignment — two roles for one geometric object.
- #disc-additive-coordinate-forcing lists Čencov-invariance as the 4th primary instance of the meta-pattern. The LMI's reliance on the FIM sits cleanly within that meta-pattern: the Fisher metric is forced by AAT-internal axioms, not adopted ad hoc.
- #result-contraction-template carries matrix-form sector machinery (Lohmiller-Slotine metric formulation) for Part III composition. The LMI's discrete-time Riccati is the control-theoretic analog of the contraction-metric framework — both are matrix-form sector conditions on multivariate dynamics.
- #def-adaptive-tempo will need a tensor extension to $\mathcal{T}$ for full alignment with this segment's directional structure; the LMI gives an explicit demand for tensor-valued adaptive rates.

Two parallel exploration drives — matrix lift. The scalar dual-drive distinction in `#disc-ci-y-unified-objective` ($\lambda_{\text{info}} \propto U_M$ epistemic vs $\lambda_{\text{surv}} \propto 1/U_M$ Lyapunov-survival) generalizes to two matrix-valued drives that compose additively:

- Λ_{info} : epistemic information-gain matrix scaled by P_t . Dominates eigendirections where P_t is large (high uncertainty in some direction) — agent explores to learn that direction.
- Λ_{surv} : Lyapunov survival-imperative matrix derived above, with eigenvalue support where the LMI constraint binds. Dominates when P_t is small along a drifting eigendirection — agent explores to track the drift it’s about to lose grip on.

The agent’s total exploration shadow price is $\Lambda_{\text{info}} + \Lambda_{\text{surv}}$, with the relative direction-by-direction weights set by where the agent sits in (P_t, Q_ρ) space.

Connection to Active Inference. Same dark-room sidestep as the scalar version: the matrix-form survival imperative is forced by Lyapunov bounds, not preferences-as-priors. The matrix lift sharpens the argument by ruling out the non-exploratory “blank wall” trivial solutions that the scalar form admitted.

Empirical signature. A 2D drift environment with separable drifting/non-drifting subspaces is the simplest test: a scalar-Lagrangian agent succumbs to blank-wall traps; an LMI agent — with Λ developing weight specifically in the drifting direction — must select actions whose J_o aligns with the drift. The directional discrimination is empirically visible in the agent’s action-selection histogram across the two subspaces.

Findings.

Matrix Lift of the Survival-Imperative Constraint via Fisher-Information LMI. **Brief:** Imagine an agent that’s supposed to track something changing but only checks for new information by staring at a blank wall. A scalar version of the survival argument would let this work — the agent satisfies “I gathered enough information” without gathering information about the thing that’s actually moving. This result fixes the math by making the argument directional: the agent has to gather information in the *directions* where the world is changing, not just gather some-information-anywhere. The scalar Causal-IB survival-imperative drive lifts to a Linear Matrix Inequality on the Fisher Information Matrix, with a positive-semidefinite matrix Lagrange multiplier Λ that distinguishes by direction. Complementary slackness mathematically forbids “blank wall” actions that satisfy the scalar bound by sourcing information in non-drifting subspaces. The fix uses Fisher-geometric machinery already AAT-internal via `#deriv-fisher-whitened-update-rule` and `#disc-additive-coordinate-forcing`’s 4th instance.

Impact: Closes a structural failure of the scalar derivation that admitted trivially-satisfying-the-math actions (constant signals, walls, etc.) without actually probing the drifting modes the agent must track to survive. The matrix lift demonstrates that the same Fisher geometric object plays multiple structurally-distinct roles — credit assignment via `#deriv-fisher-whitened-update-rule`, coordinate-forcing via `#disc-additive-coordinate-forcing`, and now survival via the LMI here. The “Tragedy of the Confident Agent” insight survives the lift in sharper form: the matrix shadow price’s eigenstructure picks out the directions where confidence has outpaced incoming information, mathematically forcing the agent to probe specifically those directions. The scalar reduction recovers `#deriv-causal-ib-exploration`’s constraint as the 1D special case, establishing structural consistency.

Novelty Claim: *Claim differentiation* on the directional discrimination of the survival-imperative exploration drive. The scalar finding in `#deriv-causal-ib-exploration` differentiated AAT’s structural-stability source for exploration from active-inference’s epistemic-value source; this finding sharpens that differentiation by ruling out the trivial-exploration solutions that the scalar form admitted, demonstrating that the structural-stability source is genuinely directional rather than scalar. The DARE-imported $J_{\min}$ and the matrix-Lagrangian KKT machinery are standard control-theoretic and convex-optimization apparatus; AAT’s contribution at this layer is the framing that places $J_{\min}$ as a survival-imperative Fisher-information floor and Λ as the directional exploration-bonus shadow price.

Related Work:

Table AV.1

<i>ASF concern</i>	<i>Prior-art language</i>	<i>Relationship / Positioning</i>
LMI / matrix-inequality formulation of stability constraints	Boyd, El Ghaoui, Feron & Balakrishnan 1994, <i>Linear Matrix Inequalities in System and Control Theory</i> , SIAM (published 1994, found 2025)	<i>formal antecedent</i> — adopted machinery; the LMI form, KKT for matrix-inequality-constrained optimization, and the DARE-imported stabilizability condition are standard control-theoretic apparatus
Discrete-time algebraic Riccati equation for steady-state Kalman survival	Anderson & Moore 1979, <i>Optimal Filtering</i> , Prentice-Hall (published 1979, found 2025)	<i>formal antecedent</i> — DARE provides the closed form for $J_{\min}(Q_P, A, R^2)$
Fisher Information Matrix as the natural geometric object for inference	Čencov 1982, <i>Statistical Decision Rules and Optimal Inference</i> (published 1982, found 2024)	<i>formal antecedent</i> — Čencov uniqueness forces the Fisher-Rao metric under (PI) parameterization-invariance, integrated AAT-internally via #disc-additive-coordinate-forcing's 4th instance
Exploration as epistemic-value drive (alternative motivational source)	Friston, Rigoli, Ognibene et al. 2015, “Active inference and epistemic value” <i>Cognitive Neuroscience</i> 6:187–214 (published 2015, found 2026-04 via Undermind Pillar 1 report)	<i>adjacent literature with structural contrast</i> — the matrix-form survival imperative is forced by Lyapunov bounds, not preferences-as-priors. The matrix lift sharpens the contrast by ruling out the non-exploratory blank-wall solutions that the scalar form admitted
Information-form Kalman filter / additive FIM update	Standard Kalman filter literature; Anderson & Moore 1979 (published 1979, found 2025)	<i>formal antecedent</i> — the additivity $P_{t+1 t+1}^{-1} = P_{t+1 t}^{-1} + \mathcal{I}_o(a_t)$ used in the LMI derivation is the information-form update step
Scalar survival-imperative drive being lifted	#deriv-causal-ib-exploration (AAT-internal, 2025)	<i>upstream segment being matrix-lifted</i> — the LMI repair preserves the dual-drive structure in matrix form ($\Lambda_{\text{info}} + \Lambda_{\text{surv}}$) and resolves the blank-wall structural failure flagged in the scalar segment's Discussion

Search Log:

- 2026-04 (*nominally comprehensive on the exploration-source contrast*, via ref/Novelty_defense_and_integration.md Pillar 1): The Undermind report on Pillar 1 (“Causal insufficiency and forced exploration”) established the active-inference EFE framing as the canonical alternative source for exploration drives. This finding's matrix-Lagrangian LMI is the directional sharpening of the scalar derivation in #deriv-causal-ib-exploration and inherits that segment's differentiation posture against EFE-style epistemic-value accounts.
- 2025 (*targeted*): Boyd et al. 1994 and Anderson & Moore 1979 confirmed as the formal antecedents for the LMI / DARE machinery; Čencov 1982 confirmed as the formal antecedent for the FIM-as-natural-geometric-object move via the (PI) axiom. AAT's contribution at this layer is framing rather than apparatus.
- 2025 (*intuition-only*): Pre-search expectation was that the LMI form and DARE machinery were standard and would not be novel apparatus; this was confirmed. The novel content is the placement — using the FIM as a survival floor rather than a controller-design parameter, and integrating directional exploration with the broader Fisher-geometric meta-pattern.

Working Notes.

- **Tensor adaptive tempo.** #def-adaptive-tempo now carries a Tensor extension sub-block ($\mathcal{T} = \sum_k v^{(k)} K^{(k)}$ with $K^{(k)} = (H_M + H_L^{(k)})^{-1} H_L^{(k)}$ from #deriv-fisher-local-update-gain) that gives the per-direction primitive this segment's LMI ma-

chinery needs. Promotion of the scalar persistence and adversarial-tempo results to invoke the tensor form directly (rather than per-coordinate scalars) remains queued — flagged in `TODD.md` under follow-on cycle work; the per-direction primitive is in place even if downstream summary results still read in scalar.

- **Worked 2D example.** A concrete 2D drift environment showing $\text{Tr}(\Lambda \mathcal{J}_o)$ distinguishing wall-channel from drift-channel actions would sharpen the blank-wall resolution. Queued.
- **2D simulation.** Generalization of `empirica/track-b-nonlinear/variants/variant_causal_ib.py` to 2D with separable drifting/non-drifting subspaces would empirically validate the matrix Lagrangian’s directional discrimination. Queued.
- **EIG vs FIM gap.** For non-Gaussian or non-local observation models, FIM diverges from EIG. The CIY-vs-EIG concern in `#disc-ci-y-unified-objective` Epistemic Status persists; the matrix lift removes the CIY- U_o surrogate but not the CIY-EIG surrogate.
- **Reasoning trail.** `spikes/.integrated/spike-causal-ib-lmi-repair.md` records the spike that surfaced the LMI repair from the scalar segment’s blank-wall critique.
- **Cross-reference to NeurIPS Paper 1.** The matrix-LMI lift is formalized at theorem grade in NeurIPS 2026 Paper 1 (“Tragedy of the Confident Agent: Forced Exploration for Survival in Drifting Environments”, `~/src/neurips/01-tragedy-confident-agent/`). The paper carries: (i) **Per-action LMI strengthenings** at §App-A.6 (`#thm-per-action-lmi`, `#thm-per-action-lmi-iid`) — pathwise survival under arbitrary action schedules with parallel high-probability bound via matrix-Chernoff concentration (`#lem-fim-concentration`), closing the pathwise-vs-expectation gap that this segment’s steady-state DARE form does not surface explicitly. (ii) **F-A-G-P enforcement framework** at §App-A.7 (`#prop-fagp`) — strict-Feasibility / drift-Aligned bonus / Gain dominance / argmax-Presentable, the four-condition operational decomposition for what makes a survival-margin controller a valid family member; the range-containment hypothesis $\text{range}(\Lambda) \subseteq \text{range}(\mathcal{J}_{\min})$ that denies blank-wall updates LMI-survival reward is the structural anchor. (iii) **Anisotropic-drift refinement** at §App-C — closed-form critical exponent p_{crit} for the $\Lambda \propto \mathcal{J}_{\min}^p$ rule, diverging as $\sigma_x/\sigma_y \rightarrow 1$; directional discrimination *rests on* anisotropic drift and the argmax flips at the isotropic limit. The scalar-isotropic case is the boundary where the matrix-LMI’s distinguishing power degrades to the scalar form. (iv) **2D Gaussian empirical anchor** with the `wall_extreme` ε -blank-wall action and the three-controller comparison (greedy / scalar / LMI) — the matrix-LMI controller selects drift-aligned probing actions and survives at 100% across drift levels at zero pragmatic reward (a *bistability* of the discrete action set, explicitly characterized in §App-B to prevent overclaiming as a smooth survival-vs-reward Pareto frontier) where scalar surrogates collapse to 0% via blank-wall degenerate updates. See `spikes/neurips-back-integration-2026-05-08.md` §1 Paper 1 entries 2, 3, 5, 6, 7 for the cross-mapping.

Appendix AW

Observation: The Authored-World Laboratory — Identifiability by Construction

AAT has one calibration laboratory it *found* — software, whose epistemic privileges are catalogued in #obs-software-epistemic-properties — and one it is *building*: a deterministic authored world (working name **vivarium**) in which the environment's law-content, state, noise, and disturbance rate are not merely observable but authored, so that every quantity the framework's results quantify over (ρ , U_o , δ_t , $\mathcal{T}$, per-dimension mismatch, $\|\delta_{\text{critical}}\|$) has ground truth at scale inside dynamics rich enough that toy-model simplifications are structurally unavailable. This segment is the canonical introduction of the authored-world laboratory, placed ahead of the framework's validation simulations because the standing intent is to re-run those simulations in it; later segments and simulation reports cite this segment rather than re-describing the system.

Formal Expression.

The authored world provides four structural properties, each the *built* counterpart of a *found* software property in #obs-software-epistemic-properties:

Empirical Claim (system properties, by construction; build-state caveat in Epistemic Status).

V1. Ground truth for every framework quantity. Law-content is an explicit versioned registry of law-articles (keyed, memoizable invocation units); state is canonical and content-addressed; the disturbance rate ρ is dialable per region (weather volatility, water-regime schedules); δ_{critical} has literal in-world meaning. Where software's inspectability is limited by the agent-side comprehension bandwidth (P1 there), the authored world's is limited only by the author's choice of instrumentation: identifiability holds **by construction** rather than by fortunate configuration.

V2. Authored coordinatization of the exogenous space. All stochasticity is a stateless pure function of a world seed and a coordinate/identity key — fated noise: genuine chance to any inside agent, deterministic lookup to the author. The aleatoric boundary's frame-relativity is thereby operational rather than philosophical, and the system realizes the central object of 7.7 — an explicit, held coordinatization of U — which makes **latent-anchored mechanism counterfactuals executable**: swap a law-article version, keep seed and keys, rerun. The separation theorem's demonstration experiment (a keyed law versus its key-permuted twin: indistinguishable at every hierarchy level, divergent under same-seed law-swap, with measurable disagreement mass) is runnable here and, at present, nowhere else available to the framework.

V3. Exact citation semantics. A claim validated in the authored world cites a specific world artifact — seed, generator versions, phase memo, intervention script — and content-addressed storage makes the citation exact and replayable to bit-identity. This is the conjunction (α) deterministic outcomes, (β) replay at commensurate cost, (γ) content-addressed immutable environment state, which #obs-software-epistemic-properties P2 identifies as jointly held, in current standard practice, by software alone — and names as a configurational uniqueness with exactly this falsifier: a simulator brought under content-addressed recording with commensurate replay cost. The authored world is that falsifier under deliberate construction; when its fidelity program completes, the “unique” in P2 becomes “one of two,” by design.

V4. A disturbance-rate schedule with a persistence reading. The world is built in phases, each running to convergence with its converged output promoted into fixed law for everything faster above it. The phase ladder is therefore a ρ -schedule: successive phases deliver environments of decreasing effective disturbance until agents of achievable tempo satisfy the persistence condition $\mathcal{T} > \rho / \|\delta_{\text{critical}}\|$ (4.5). An agent-bearing world is, in this reading, the terminal rung of a persistence ladder the authors descend deliberately.

Epistemic Status.

Discussion-grade as a whole, with the property claims empirical-by-construction and honestly dated. V1-V4 describe the system’s design and its verified build state as of 2026-07-16: the world is mid-Phase-3 (geology and hydrology converged with absolute-tolerance conservation tests; climate next; **agents not yet embodied**), so V1-V3 are live properties of the running system while V4’s agent-bearing terminus and everything downstream of the agent seam are design intent. Claims here about what the laboratory *will* offer AAT’s agent-level results are forward-looking and marked so; claims about determinism, keyed noise, content-addressing, and conservation-honesty describe running code. The Level-4 executability claim (V2) inherits *exact* semantics from 7.7; that the system realizes the required coordinatization is a fact of its architecture. This segment deliberately does not depend on the world’s agent layer existing: it introduces the laboratory, not results obtained in it.

Discussion.

Found versus built, and why the framework wants both. Software is the domain where the epistemic privileges happen to hold and a large practice already generates data inside them; the authored world is the domain where they hold because someone chose them, which buys three things software cannot: (a) *dialable* environment parameters — ρ can be scheduled, not merely measured; (b) *author-side access to the exogenous coordinatization* — software’s version control gives literal Level-3 access (P2 there), while fated noise gives literal Level-4 access, one full rung more; (c) freedom from the confound that software environments are themselves made of agent behavior. Conversely the authored world lacks software’s ecological validity — nothing in it ships to production — so results transfer with stated construction assumptions, exactly the transfer-assumption discipline TST names for cross-domain claims. The two laboratories are complements, not competitors: found breadth against built depth.

The standing simulation intent. The framework’s Part I validation simulations (AX) were built as standalone numerical models; the standing intent is to re-run them in-vivo in the authored world as its phases reach the relevant dynamics, converting each from a bespoke script into a cited, replayable world artifact per V3. This segment exists partly so that migration has a canonical anchor: simulation reports should cite the laboratory, not re-describe it.

Near-term experiments the laboratory makes cheap. Beyond the Level-4 demonstration (V2), two are flagged in adjacent segments: the GA-1 housing experiment (transition-law versus observation-law ignorance — the experimenter chooses whether to perturb T or h , directly testing the verified asymmetry from the epistemic-target work), and the behavioral $\hat{\kappa}_{\text{processing}}$ estimator at scale (same event under different goal states across thousands of agent-events with controlled priors) once the agent seam lands.

Ethics gate. Everything agent-seam-ward is gated by the authored world's own ethics front door, which is deliberately more conservative than this framework's volume-04 gates and binds program-wide; the laboratory's agent-level offers are sequenced behind it, not around it.

Working Notes.

- **Provenance and living bridge.** Canonical intro landed 2026-07-16 from the vivarium handshake work (same-day: the mechanism-counterfactual mapping). The living cross-project bridge — current build state, the four-thread Level-4 handshake, near/far offers — is `doc/vivarium.md` (repo-internal working doc, not citable from canon bodies; this WN pointer is the sanctioned breadcrumb). Vivarium-side twin: `~/src/archema-io/vivarium/feedback-from-asf.md` §Level-4 handshake; its ethics front door is `~/src/archema-io/vivarium/ETHICS.md`.
 - **Build-state re-verification.** The Epistemic Status's build-state sentence is dated 2026-07-16; re-verify against the vivarium tree before citing this segment's V-claims in any promoted result (mid-rebuild volatility is a known divergence zone per the program's incoherence ledger).
 - **Chronica link (forward, cross-volume).** A world's interventional history is a chronica in the exact sense of 04-eli-core's `#def-chronica` (ordinal, causal, append-only), and the world's fork/save machinery is the operational subject of the proposed `#hyp-checkpoint-forking-failure-modes` — the two should co-develop; kept as a WN forward pointer rather than a depends edge to preserve volume reading order.
 - **P2 uniqueness-claim maintenance.** When the fidelity program completes (γ at full strength), `#obs-software-epistemic-properties` P2's "unique in current standard practice" needs its by-then-true revision; a matching WN pointer lives there.
-

Appendix AX

Detail: Observation: Part I Validation Simulations

Six simulation variants validated, refined, and extended the analytical predictions from Part I, forcing regime splits and discovering scaling laws that the analytical derivations left ambiguous.

The simulations are **theory-shaping, not merely confirmatory** — three findings emerged from simulation that *forced refinements* to the analytical treatment: (1) the adversarial tempo exponent turned out to be regime-dependent (Model D vs Model S; coupling-dominant vs non-coupling-dominant), forcing the $b = 2 / b = 3/2 / b \rightarrow 1$ regime classification the linear ODE left ambiguous; (2) observation noise was found to gate adversarial advantage with optimal-gain partial recovery, a result that had no analytical treatment before the simulations; (3) per-dimension persistence emerged as load-bearing — anisotropic systems with gain varying across dimensions produced mismatch concentration on the weak direction that the scalar persistence condition systematically overestimated by up to 72%. The simulation program is treated as part of AAT's methodology rather than as external validation: where simulation surfaces something the analytical treatment misses, the analytical treatment is refined; where simulation confirms what the analytical treatment predicts, the result lifts in confidence.

Overview.

The track-b simulation program tested Part I's analytical predictions about mismatch dynamics, adversarial coupling, observation noise, correction function robustness, and anisotropic persistence. The simulations modeled discrete-time mismatch dynamics as AR(1) processes with five correction functions (linear, saturating, threshold/dead-zone, sigmoid, structural breakdown), sweeping parameter spaces across gain, disturbance rate, coupling strength, observation noise, and dimensional structure.

The simulations were theory-shaping, not merely confirmatory. In particular:

- The adversarial tempo exponent turned out to be regime-dependent (deterministic drift vs. stochastic noise, coupling-dominant vs. non-coupling-dominant), forcing a regime split that the mismatch ODE left ambiguous.
- Observation noise was found to gate the adversarial advantage, partially recoverable by optimal gain – a finding that had no analytical treatment before the simulations.
- The scalar persistence condition was shown to overestimate adaptive capacity by 72% in anisotropic systems, motivating the per-dimension persistence condition.

Variant Summary.

Table AX.1

<i>Variant</i>	<i>Disturbance Model</i>	<i>Tested</i>	<i>Key Finding</i>	<i>Theory Impact</i>	<i>Promoted Segment</i>
A	D (deterministic drift)	Deterministic drift coupling; coupling-dominance sweep	Exponent $b \rightarrow 2.0$ in coupling-dominant limit (confirmed at 1.999)	Validates Model D derivation ($b = 2$)	16.3
B	D+S (interpolation)	Drift-noise interpolation across correction functions	Smooth transition between drift ($b = 2.0$) and noise ($b = 1.5$) regimes; coupling dominance is the key qualifier, not drift vs. noise per se	Confirmed that the coupling-dominant condition is quantitatively load-bearing	16.3
C	S (stochastic AR(1))	Exponent vs. gain (η) in stochastic model	Exponent drops toward 0.5 as $\eta \rightarrow 0$ (away from coupling dominance at fixed q_{base}); discrete AR(1) exponent never exceeds 1.5	Validates Model S derivation ($b = 3/2$)	16.3
D	D vs. S (separated)	Exponent vs. base noise (q_{base}) at fixed η	Continuous ODE exponent $\rightarrow 2.0$ (Model D), discrete AR(1) exponent $\rightarrow 1.5$ (Model S), as $q_{\text{base}} \rightarrow 0$	Definitively separated the two asymptotes; validates both derivations	16.3
E	S (stochastic AR(1))	Observation noise; optimal gain validation	Observation noise collapses adversarial exponent from ~ 1.0 to ~ 0.2 ; Riccati-optimal gain restores it to ~ 0.4 ; 52% mismatch reduction at moderate noise	Observation quality gates tempo advantage; optimal gain (3.6) empirically validated	16.4
F	S (stochastic AR(1))	Multi-dimensional anisotropic correction; targeted adversarial attack	Per-dimension theory exact to 4 significant figures; scalar tempo overestimates by 72%; targeted attack amplifies advantage by 17%	Validates Model S per-dimension steady state	16.5
Hafez bridge	S (stochastic AR(1))	Bi-predictability P vs. AAT mismatch in adversarial and non-adversarial settings	P measures coupling architecture (scale-invariant); mismatch measures coupling performance; P is blind to adversarial dynamics	P and AAT mismatch are complementary diagnostics; H_b (agent opacity) has no direct AAT analog – potential gap for multi-agent work	–

Methodology.

Discrete mismatch dynamics. The environment follows a random walk $x_{t+1} = x_t + w_t$ with $w_t \sim N(0, q^2)$. The agent corrects via $\hat{x}_{t+1} = \hat{x}_t + \eta \cdot g(\delta_t)$, yielding the AR(1) mismatch process $\delta_{t+1} = (1 - \eta) \cdot \delta_t + w_t$ for linear g . This is the discrete-time analog of the mismatch ODE $d\|\delta\|/dt = -\mathcal{T} \cdot \|\delta\| + \rho$ from 3.9.

Parameter sweeps. Each variant swept its key parameter(s) across 7–20 values. Monte Carlo: 200 independent trials per parameter point, 10,000–20,000 timesteps per trial, with 2,000–5,000 step burn-in for steady-state convergence. Fixed random seeds for reproducibility.

Exponent fitting. Adversarial exponents were estimated by fitting $\log(\|\delta_B\|/\|\delta_A\|) = a + b \cdot \log(\mathcal{T}_A/\mathcal{T}_B)$ via weighted least squares, with weights inversely proportional to variance of each point's log-estimate. 95% confidence intervals via bootstrap (1,000 samples).

Correction functions. Five functions $g : \mathbb{R} \rightarrow \mathbb{R}$ were tested, all satisfying $g(0) = 0$ and $g'(0) = 1$: linear ($g(\delta) = \delta$), saturating ($g(\delta) = \delta/(1 + |\delta|/R)$), threshold ($g(\delta) = \delta \cdot \mathbf{1}[|\delta| > \epsilon]$), sigmoid ($g(\delta) = R \cdot \tanh(\delta/R)$), and structural breakdown ($g(\delta) = \delta \cdot \mathbf{1}[|\delta| < R_{\max}]$).

Simulation code (reproducibility). The full simulation suite — single-agent and two-agent initial simulations plus the four variant-extension drivers (A/B drift, C/D regimes, E/F extensions, Hafez bridge) and their per-variant result write-ups — is maintained as the project's reproducibility supplement and will be released as a citable code artifact at publication. The numerical outcomes that the body's findings rest on (steady-state formulas, the fitted exponents, the gating thresholds, the per-dimension errors) are stated self-contained in the tables and Key Findings below.

Key Findings.

The adversarial exponent is regime-dependent (now derived). The theory now distinguishes two disturbance models (Model D: bounded deterministic, GA-2; Model S: stochastic zero-mean, GA-2S), each producing a different steady-state scaling and therefore a different adversarial exponent:

- **Model D, coupling-dominant:** $\|\delta\|_{ss} = \rho/\mathcal{T}$. Adversarial exponent $b = 2$ (derived from Prop A.1; Variant A confirms at 1.999).
- **Model S, coupling-dominant:** $\|\delta\|_{rms} = \sigma_w/\sqrt{2\mathcal{T}}$ (from Prop A.1S). Adversarial exponent $b = 3/2$ (derived; Variants C-D confirm at 1.481).
- **Non-coupling-dominant:** Exponent degrades smoothly toward 1.0 (Model D) or 0.5 (Model S) as base disturbance grows relative to adversarial coupling.

A measured exponent of $b \approx 1.05$ corresponds to the non-coupling-dominant regime (stochastic-noise coupling at moderate coupling dominance), consistent with the regime classification rather than contradicting it. Variants A–D systematically map the full regime space and validate the derived exponents. See 16.3 for the full treatment.

Observation noise gates adversarial advantage. Adding observation noise σ_{obs} to the mismatch signal collapsed the adversarial exponent from ~ 1.0 to ~ 0.2 at $\sigma_{obs} = 10 \times q_{env}$ (fixed gain). The Riccati-optimal gain (3.6) partially restored the advantage to ~ 0.4 , more than doubling it but not recovering the noise-free level. The optimal gain helps most in the moderate-noise regime, where it achieved a 52% mismatch reduction over fixed gain. See 16.4.

Per-dimension persistence is exact; scalar overestimates. In a 3-dimensional system with 5:1 gain variation across dimensions, the per-dimension AR(1) steady-state prediction matched simulation to 4 significant figures. The scalar persistence condition overestimated aggregate mismatch by 72%, because it averages across dimensions while the weak dimension dominates the L_2 norm. Isotropic gain allocation reduced overall mismatch by 13%; targeted adversarial attack on the weak dimension amplified the mismatch ratio by 17%. See 16.5.

Nonlinear correction creates thresholds, not lower exponents. Under deterministic drift, saturating, sigmoid, and breakdown correction functions did not simply reduce the adversarial exponent. Instead, they produced catastrophic divergence when disturbance exceeded the correction capacity ($\rho > \mathcal{T} \cdot R$). This is the persistence threshold failure from 4.5, observed directly in simulation. The measured “exponents” above 3.0 for these functions were divergence artifacts, not meaningful scaling laws.

Hafez bridge: architecture vs. performance. Bi-predictability P (Hafez et al.) measures the informational architecture of agent-environment coupling; AAT mismatch measures the operational performance. In adversarial settings, P remained nearly constant (~ 0.268) across a 10:1 range of opponent tempo, while the mismatch ratio varied from 0.54 to 5.95. P is scale-invariant after z-score normalization, making it insensitive to adversarial dynamics. The two metrics are complementary: P diagnoses whether the coupling structure is sound; mismatch predicts how well it performs.

Epistemic Status.

Empirical. Simulation results are reproducible under fixed random seeds; the driver code and per-variant write-ups are maintained as a reproducibility supplement to be released as a citable code artifact at publication. The key results – regime-dependent exponents, observation noise gating, per-dimension exactness – have been promoted to first-class segments (16.3, 16.4, 16.5) with their own epistemic assessments. This appendix serves as reference for the simulation program as a whole.

Discussion.

Internal validation, not external. The simulations operate in the model’s own terms: AR(1) processes with parameterized correction functions, Gaussian noise, and discrete-time dynamics. They validate the mathematical predictions within the model – confirming that the analytical steady-state formulas, the exponent regimes, and the per-dimension theory are correct as statements about these stochastic processes. The gap between “the math predicts X within its own model” and “real agents exhibit X” remains an empirical question for each domain.

What the simulations did not test. The simulations did not test whether the AR(1) mismatch dynamics are a good model for any particular real-world adaptive system. They also did not test cross-dimensional coupling (off-diagonal correction), non-Gaussian disturbances, non-stationary parameters, or multi-agent composition dynamics. The interaction between observation noise and the deterministic-drift exponent regime ($b = 2.0$) was not tested – Variant E used stochastic coupling only.

Theory-shaping role. The most important contribution of the simulations was not confirming predictions but forcing the theory to be more precise. The regime split in the adversarial exponent, the observation-noise gating mechanism, and the per-dimension bottleneck effect were all findings that the analytical derivations initially left ambiguous or unaddressed. The simulations forced the disturbance model split (Model D vs. Model S), which in turn enabled the analytical derivation of both exponents. The simulations now serve as validation of the derived results: Variant A validates $b = 2$ (Model D), Variants C-D validate $b = 3/2$ (Model S), Variant F validates the Model S per-dimension steady state.

Part V

Appendices: Operational Domains

Appendix AY

Detail: Operationalization — Estimation Procedures

Estimation recipes for core AAT quantities, bridging the measurement gap between formal objects and practical deployment.

The segment is **domain-honest** about which quantities are easily measured and which are operationally challenging: model uncertainty admits predictive-posterior / ensemble-spread / Bayesian-NN estimators; observation uncertainty admits channel calibration or residual variance; mismatch magnitude is directly measurable from prediction-versus-observation residuals; tempo from event rate times effective gain; the sector parameter via either the gain-sector bridge (where applicable — Bayesian on exponential families; gradient on strongly convex) or empirical correction-rate-vs-mismatch regression; disturbance parameters split per Model D ($\hat{\rho}_{\text{det}}$) versus Model S ($\hat{\sigma}_w$) with sequential dependence on $\hat{\mathcal{T}}$; critical-mismatch tolerance is a *domain parameter* not an estimated quantity (it encodes “how wrong can the model be before actions become harmful or ineffective?”). Model-class fitness is named as particularly challenging — not directly computable, requiring search across the model class; the observable signature is persistent mismatch despite adequate learning. Satisfaction gap and control regret require evaluating attainability, generally intractable in exact form. AAT’s commitment is that the abstract machinery is operationalizable with stated quality trade-offs at each estimation step, not analytical idealization with no path to practice.

Measurement Targets.

Table AY.1

<i>Quantity</i>	<i>Role in AAT</i>	<i>Typical unit</i>	<i>Minimum data needed</i>
U_M	Model uncertainty (3.6)	domain-specific variance/entropy	Predictive posterior or ensemble spread
U_o	Observation uncertainty (3.6)	sensor variance/noise scale	Channel calibration or residual variance
$\rho_{\text{det}}(t)$	Mismatch injection rate bound, Model D (3.9)	surprise per time	Time-series of mismatch magnitudes; worst-case bounds
σ_w	Disturbance noise scale, Model S (3.9)	surprise per $\sqrt{\text{time}}$	Mismatch innovation variance (residual after correction)
ρ_{delib}	Local mismatch drift during pauses (4.1)	surprise per time	Deliberation windows with no corrective action
α	Lower correction efficiency bound (4.4)	inverse time	Vector mismatch trajectories + correction term

<i>Quantity</i>	<i>Role in AAT</i>	<i>Typical unit</i>	<i>Minimum data needed</i>
R	Radius where local sector condition holds (4.4)	surprise magnitude	Same as α , plus breakdown detection
$\ \delta_{\text{critical}}\ $	Functional adequacy threshold (4.5)	surprise magnitude	Task-level performance curve vs mismatch

Estimator Cookbook.

Estimating U_M and U_o . Use the most native uncertainty representation available in the domain:

Table AY.2

<i>Domain</i>	<i>U_M estimator</i>	<i>U_o estimator</i>
Kalman / linear Gaussian	Prior predictive variance $P_{t t-1}$	Measurement-noise variance R_t (known or EM-estimated)
Conjugate Bayes	Posterior variance / inverse effective sample size	Likelihood variance (or precision inverse)
RL with ensembles	Across-head predictive variance $\text{Var}_t[Q_t(s, a)]$	TD-target noise variance over replay batches
Neural net regression	Ensemble or Laplace posterior variance	Aleatoric head output $\sigma^2(x)$
PID / classical control	State-estimation covariance from observer	Sensor noise from calibration + residual PSD

For the scalar gain heuristic (3.6), normalize to common units and compute:

Operational Definition.

$$\hat{\eta}_t^* = \frac{\hat{U}_{M,t}}{\hat{U}_{M,t} + \hat{U}_{o,t}} \quad (\text{AY.1})$$

Estimating Disturbance Parameters (ρ_{det} , σ_w , ρ_{delib}). The disturbance model split (Model D vs. Model S; see 3.9, 4.4) requires different estimation procedures depending on the disturbance character.

Identifying the disturbance model. The choice between Model D and Model S is domain-dependent: - **Model D** (bounded deterministic, GA-2): appropriate when the environment changes persistently and directionally. Estimate ρ_{det} from worst-case bounds on mismatch injection rate. Examples: adversary who maneuvers systematically, API that drifts, climate that shifts. - **Model S** (stochastic zero-mean, GA-2S): appropriate when the environment fluctuates unpredictably around a stable mean. Estimate σ_w from residual variance of mismatch innovations. Examples: market noise, sensor noise, random perturbations. - **Mixed environments** use both: estimate ρ_{det} from the persistent component (trend) and σ_w from the residual variance after detrending. The adversarial exponent interpolates smoothly between $b = 2$ (pure drift) and $b = 3/2$ (pure noise); see 16.3.

Let $s_t = \|\delta_t\|$ in surprise units (e.g., negative log-likelihood residual scale).

Model D: Global mismatch injection rate ($\hat{\rho}_{\text{det}}$):

Operational Definition.

$$\hat{\rho}_{\text{det}}(t) = \left[\frac{s_{t+\Delta t} - s_t}{\Delta t} + \hat{\mathcal{F}}_t s_t \right]_+ \quad (\text{AY.2})$$

where $[x]_+ = \max(x, 0)$ and $\hat{\mathcal{F}}$ evaluated at time t is estimated adaptive tempo. Take the supremum (or a high quantile) over the estimation window.

Model S: Noise scale ($\hat{\sigma}_w$):

Operational Definition.

$$\hat{\sigma}_w^2 = \text{Var} \left[\frac{s_{t+\Delta t} - (1 - \hat{\eta}^*)s_t}{\sqrt{\Delta t}} \right] \quad (\text{AY.3})$$

i.e., the variance of the mismatch innovations (residuals after subtracting the expected correction). For AR(1) processes, this is the innovation variance $\hat{\rho}^2$ directly.

Note on estimation sequencing. Both estimators require $\hat{\mathcal{F}}_t$ (or $\hat{\eta}^*$), estimated from $\hat{\nu}$ and $\hat{\eta}^*$. Estimate the gain and event rate first (from the agent's internal statistics and observation timing), then use these to extract the disturbance parameters from the mismatch trajectory. This sequential structure avoids circularity but introduces sensitivity: errors in $\hat{\mathcal{F}}$ propagate into $\hat{\rho}_{\text{det}}$ and $\hat{\sigma}_w$.

Local pause-window drift for Derived 4.2 ($\hat{\rho}_{\text{delib}}$):

Operational Definition.

$$\hat{\rho}_{\text{delib}} = \text{median}_{w \in \mathcal{W}_{\text{pause}}} \frac{s_{w,\text{end}} - s_{w,\text{start}}}{\Delta \tau_w} \quad (\text{AY.4})$$

using windows where corrective action is suspended or effectively delayed.

Estimating α (sector lower bound). A uses $\delta^T F(\mathcal{T}, \delta) \geq \alpha \|\delta\|^2$ for $\|\delta\| \leq R$. Operationally:

1. Estimate $\dot{\delta}_t$ (finite differences or filtered derivative).
2. Compute $\hat{F}_t = -\dot{\delta}_t + w_t$ where disturbance proxy w_t is estimated from exogenous perturbation channels or residual balancing.
3. Form ratios $r_t = (\delta_t^T \hat{F}_t) / \|\delta_t\|^2$ on bins of $\|\delta_t\|$.
4. Set conservative lower bound $\hat{\alpha}$ as a low quantile (e.g., 10th percentile) of r_t in the valid region.

Estimating R (valid-region radius). Estimate R as the largest radius for which sector inequality violations remain below tolerance:

Operational Criterion.

$$\hat{R} = \sup \{r > 0 : \Pr(\delta^T \hat{F} < \hat{\alpha} \|\delta\|^2 \mid \|\delta\| \leq r) \leq \epsilon\} \quad (\text{AY.5})$$

with a chosen violation tolerance ϵ (e.g., 5%).

Estimating $\|\delta_{\text{critical}}\|$. Define a mission-level performance metric J (reward rate, tracking error, service SLA, etc.). Set the critical mismatch threshold where performance crosses a minimum acceptable level $J_{\min}$:

Operational Definition.

$$\|\hat{\delta}_{\text{critical}}\| = \inf \{d : \mathbb{E}[J \mid \|\delta\| = d] < J_{\min}\} \quad (\text{AY.6})$$

This anchors 4.5 to real task outcomes.

Recommended Estimation Sequence.

1. Fix mismatch representation δ in one consistent unit system (prefer surprise-scale).
2. Estimate U_o from channel physics/calibration; estimate U_M from model uncertainty.
3. Validate gain behavior against 3.6 ($\hat{\eta}^*$ trend checks).
4. Estimate ρ_{delib} from pause windows (4.1) and $\rho(t)$ from full traces.
5. Estimate α and R from local correction dynamics (A).
6. Estimate $\|\delta_{\text{critical}}\|$ from task-performance degradation.
7. Compute derived diagnostics:
 - Tempo margin (Model D): $\hat{\mathcal{T}} - \hat{\rho}_{\text{det}} / \|\hat{\delta}_{\text{critical}}\|$
 - Tempo margin (Model S): $\hat{\mathcal{T}} - n\hat{\sigma}_w^2 / (2\|\hat{\delta}_{\text{critical}}\|^2)$
 - Reserve: $\widehat{\Delta\rho^*} = \hat{\alpha}\hat{R} - \hat{\rho}_{\text{det}}$
 - Deliberation feasibility: $\Delta\eta^*(\Delta\tau)\|\delta_{\text{post}}\| - \hat{\rho}_{\text{delib}}\Delta\tau$

Decision-Theoretic Procedures.

<i>Context</i>	<i>λ estimator</i>	<i>Source</i>
----------------	---------------------------------------	---------------

Estimating λ (Exploration Price). The exploration weight $\lambda(M_t)$ in the unified policy objective (7.5) prices information in value-equivalent terms:

Table AY.3

<i>Context</i>	<i>λ estimator</i>	<i>Source</i>
Finite bandits	Gittins index from dynamic programming	Exact (Gittins 1979)
Linear-Gaussian	Probing cost in quadratic objective	Exact (dual control)
Discrete MDP	(VoI) ² /info gain	Information-directed sampling (Russo & Van Roy)
General	$\hat{\lambda} = c \cdot \hat{U}_M / \hat{U}_o$	Heuristic: scale CIY weight by relative uncertainty

For the heuristic: when $U_M \gg U_o$ (highly uncertain model), exploration is cheap relative to exploitation risk, so λ should be large. When $U_M \ll U_o$ (confident model, noisy observations), exploitation dominates. The constant c is domain-specific.

Deliberation Stopping Policy. From 4.1, deliberation of duration $\Delta\tau$ is warranted when $\Delta\eta^*(\Delta\tau) \cdot \|\delta_{\text{post}}\| > \rho_{\text{delib}} \cdot \Delta\tau$. Operationally:

1. Estimate ρ_{delib} from prior pause windows.
2. Before each deliberation episode, estimate $\|\delta_{\text{post}}\|$ as current mismatch + $\rho_{\text{delib}} \cdot \Delta\tau_{\text{planned}}$.
3. Estimate $\Delta\eta^*(\Delta\tau)$ from the diminishing-returns profile of past deliberation episodes.
4. Stop deliberating when the marginal improvement rate $\partial\Delta\eta^*/\partial\Delta\tau$ drops below $\rho_{\text{delib}}/\|\delta_{\text{post}}\|$.

Structural-Switch Trigger. From 4.8, structural adaptation is indicated when parametric convergence leaves a mismatch floor. Operationally:

1. Estimate the current mismatch floor $\|\delta\|_{\text{floor}}$ from converged residual statistics.
2. Estimate post-switch expected mismatch as $\|\delta\|_{\text{new}} \approx \rho/\alpha'$ where α' is the sector bound under the candidate new model class.
3. Estimate transition cost C_{switch} : knowledge loss, retraining time ($\Delta\tau_{\text{switch}}$), and accumulated mismatch during transition ($\rho \cdot \Delta\tau_{\text{switch}}$).
4. Switch when: $(\|\delta\|_{\text{floor}} - \|\delta\|_{\text{new}}) \cdot T_{\text{horizon}} > C_{\text{switch}}$.

Estimator Uncertainty Guidance.

$\hat{\eta}^*$ (**gain estimate**). Approximate variance via the delta method:

$$\text{Var}(\hat{\eta}^*) \approx \hat{\eta}^{*2}(1 - \hat{\eta}^*)^2 \left[\frac{\text{Var}(\hat{U}_M)}{\hat{U}_M^2} + \frac{\text{Var}(\hat{U}_o)}{\hat{U}_o^2} \right] \quad (\text{AY.7})$$

Near 0 or 1, stable (dominated by one source). Near 0.5, most volatile.

$\hat{\rho}$ (**mismatch injection rate**). Finite-difference estimation amplifies noise. Recommend smoothing before differencing. Minimum ~ 20 observations for stable trend; ~ 50 for reliable variance. The $[\cdot]_+$ clipping introduces positive bias at small ρ .

$\hat{\alpha}$ (**sector lower bound**). The 10th-percentile approach gives approximately 90% confidence under stationarity. Report quantile level, bin count, and bin width.

$\hat{R}$ (**sector-condition radius**). A sharp drop-off in sector-condition satisfaction is a strong signal; gradual decay suggests the sector condition may not hold cleanly.

Nonstationarity caveat. All estimators assume approximate stationarity over the estimation window. When environment regime changes are suspected (4.8), re-estimate from post-change data only.

Epistemic Status.

This is a *procedures document* — estimation recipes, not theoretical claims. Each estimator inherits the epistemic status of the quantity it targets: the gain estimator’s validity depends on 3.6’s structural claim; the sector-bound estimator’s validity depends on 4.4’s assumptions. The estimation procedures themselves are *conditional* on these theoretical foundations and on the stationarity and data-sufficiency assumptions noted above.

Working Notes.

- End-to-end worked examples instantiating the full chain are in [AZ](#) (exact) and [BA](#) (approximate).
 - The $\hat{\rho}$ estimator’s dependence on $\hat{\mathcal{T}}$ creates a sequential estimation chain. Errors compound. An alternative approach: estimate ρ directly from exogenous environmental change measurements when available, bypassing the mismatch trajectory entirely.
 - The structural-switch trigger’s T_{horizon} (expected time the new model class will remain adequate) is itself uncertain and difficult to estimate. In practice, this is where the agent’s M_t does the real work — predicting environmental stability.
-

Appendix AZ

Worked Example: 1D Active Tracking (Kalman Domain)

Every AAT quantity has an exact Kalman-filter counterpart. This is a *validation* of the formal chain — all quantities are computable in closed form.

This is AAT's **calibration on a problem with known exact answers**. AAT's contribution in the Kalman instance is not adding new results to Kalman theory; it is *embedding Kalman in a more general framework that handles the same primitives across non-linear, non-Gaussian, and non-fully-observable agent classes*, with the Kalman case as the special case where everything reduces to closed-form. The mismatch is the innovation; the optimal gain coincides with the Kalman gain (the uncertainty-ratio formula equals the Kalman gain in the linear-Gaussian case); the sector parameter equals tempo via 4.3; the persistence condition reduces to observability plus detectability. Where AAT predicts what Kalman theory says, the framework is minimally correct; where AAT predicts beyond Kalman's scope (non-linear, non-Gaussian, partial observability, multi-agent), the framework's reach extends past Kalman's. The Kalman instance is the anchor point that grounds the abstract machinery in a familiar concrete case.

System.

The agent tracks scalar state x_t and chooses sensor mode $a_t \in \{L, H\}$:

- Dynamics: $x_{t+1} = x_t + v_t$, $v_t \sim \mathcal{N}(0, q)$, $q = 0.25$
- Observation: $y_t = x_t + n_t^{(a_t)}$
- Low mode noise: $r_L = 9$; High mode noise: $r_H = 1$
- Event rate: $\nu = 5$ Hz
- High mode has higher energy cost

Chain Instantiation.

Scope (1.5). *Mapping: exact.*

$\Omega_t = x_t$ is partially observed via noisy channel. $\mathcal{A} = \{L, H\}$ is non-empty and causally affects observation quality. Residual uncertainty persists due to process and sensor noise.

Causal Structure (1.7) + CIY (3.7). *Mapping: exact.*

Action precedes observation and changes $P(y_t \mid do(a_t))$ through r_{a_t} . Using low mode as comparator action, the CIY is the interventional KL divergence:

Worked Quantity.

$$\text{CIY}(H) = D_{\text{KL}}(P(y \mid do(H)) \parallel P(y \mid do(L))) \quad (\text{AZ.1})$$

$$= \frac{1}{2} \left[\log \left(\frac{P^- + r_L}{P^- + r_H} \right) + \frac{P^- + r_H}{P^- + r_L} - 1 \right] \quad (\text{AZ.2})$$

With $P^- = 4.25$: $\text{CIY}(H) \approx 0.161$ nats.

Model (2.1). *Mapping: exact.*

Model state $M_t = (\hat{x}_{t|t}, P_{t|t})$ — a compression of interaction history with recursive update.

Mismatch (3.4). *Mapping: exact.*

$$\delta_t = y_t - \hat{x}_{t|t-1} \quad (\text{AZ.3})$$

Update Gain (3.6). *Mapping: exact.*

Scalar Kalman gain:

$$K_t = \frac{P_t^-}{P_t^- + r_{a_t}} \quad (\text{AZ.4})$$

With $P^- = 4.25$: $K(H) \approx 0.810$, $K(L) \approx 0.321$.

The exact uncertainty ratio mapping: $U_M = P_t^-$, $U_o = r_{a_t}$, $\eta^* = K_t$.

Exploration (3.7).

Worked Objective.

$$a_t^* = \arg \max_a [\mathbb{E}[\text{value}(a) \mid M_t] + \lambda_t \mathbb{E}[\text{CIY}(a) \mid M_t]] \quad (\text{AZ.5})$$

When uncertainty is high (P^- large), CIY term favors high mode. As uncertainty falls, policy shifts toward low-cost L .

Deliberation Threshold (4.1). Suppose a planning pause of $\Delta\tau = 0.5$ s, with measured $\rho_{\text{delib}} = 0.40$ surprise/s.

Cost during pause: 0.20 surprise units. If $\|\delta_{\text{post}}\| = 0.70$, deliberation is worthwhile when:

$$\Delta\eta^*(0.5) \cdot 0.70 > 0.20 \implies \Delta\eta^*(0.5) > 0.286 \quad (\text{AZ.6})$$

Structural Adaptation (4.8). Assume maneuvering regime change introduces sustained residual autocorrelation and mismatch floor. If estimated valid radius drops to $R = 0.08$ while $R^* = \rho/\alpha = 0.12$, parametric adaptation is no longer adequate ($R^* > R$), triggering model-class change (e.g., constant-velocity $\rightarrow$ constant-acceleration process model).

Tempo + Persistence (3.8, 4.5). Using action mix 70% H , 30% L :

$$\bar{\eta}^* = 0.7(0.810) + 0.3(0.321) = 0.663 \quad (\text{AZ.7})$$

$$\mathcal{T} = \nu \bar{\eta}^* = 5 \cdot 0.663 = 3.315 \text{ s}^{-1} \quad (\text{AZ.8})$$

With $\rho = 0.18$ surprise/s and $\|\delta_{\text{critical}}\| = 1$:

$$\mathcal{T} > \frac{\rho}{\|\delta_{\text{critical}}\|} \Rightarrow 3.315 > 0.18 \checkmark \quad (\text{AZ.9})$$

Lyapunov Bounds (4.4, 4.3). *Mapping: exact (derived, not estimated).*

By 4.3, the sector parameter for a scalar Kalman filter is $\alpha = \mathcal{T} = \nu \bar{\eta}^*$:

$$\alpha = \mathcal{T} = 3.315 \text{ s}^{-1} \quad (\text{AZ.10})$$

This is derived from the Kalman gain, not estimated from data — the correction function $F(e) = Ke$ is exactly linear, so $\alpha = K$ per observation and $\alpha = \nu \bar{K} = \mathcal{T}$ in continuous time. With $R = 1.4$ (model class capacity, a domain parameter) and $\rho = 0.18$:

$$R^* = \frac{\rho}{\alpha} = \frac{0.18}{3.315} \approx 0.054 < R \quad (\text{AZ.11})$$

$$\Delta\rho^* = \alpha R - \rho = 3.315(1.4) - 0.18 = 4.46 \quad (\text{AZ.12})$$

The agent is comfortably within its invariant region with substantial adaptive reserve.

Mapping Quality Summary.

Table AZ.1

<i>AAT Concept</i>	<i>Kalman Mapping</i>	<i>Status</i>
Scope	Exact	Definitional
Causal structure + CIY	Exact	Closed-form KL
Model (M_t as sufficient statistic)	Exact	Kalman state + covariance
Mismatch ($\delta_t = \text{innovation}$)	Exact	Standard Kalman innovation
Gain ($\eta^* = K_t$)	Exact	Kalman gain IS uncertainty ratio
Tempo ($\mathcal{T} = \nu \bar{\eta}^*$)	Exact	Closed-form
Persistence condition	Exact	Linear ODE solution
Sector parameter ($\alpha = \mathcal{T}$)	Exact	Derived from Kalman gain via 4.3

<i>AAT Concept</i>	<i>Kalman Mapping</i>	<i>Status</i>
Lyapunov bounds ($R^*, \Delta \rho^*$)	Exact	From derived sector parameter + domain R

Epistemic Status.

This is a *worked instantiation*, not a theoretical claim. Every mapping is exact — the Kalman domain is the canonical case where AAT’s formal chain has closed-form realizations. The example validates that the formal chain is internally consistent and instantiable.

Working Notes.

- This example uses a 1D system. The tensor-tempo extension (16.5) becomes visible only in multi-dimensional tracking where different state dimensions have different observability.
- The structural adaptation trigger (constant-velocity $\rightarrow$ constant-acceleration) is manufactured for the example. A more natural test would be a real tracking system with genuine regime changes.
- The Kalman filter is provably optimal for the linear-Gaussian case. Every AAT quantity has not just an analog but the *exact optimal* value. This makes it the strongest validation but also the easiest — the real test is non-Kalman domains (see BA).

Appendix BA

Worked Example: Nonstationary Bandit (RL Domain)

AAT’s conceptual architecture applies beyond the conjugate-Bayesian regime. The mapping here is *approximate* — it shows that AAT’s concepts organize RL phenomena, but the quantitative relationships are structural analogies, not derivations. Where the mapping is exact versus approximate is marked explicitly.

This is the framework’s bridge to RL theory. AAT does not subsume RL theory; it claims that typical RL methods (ϵ -greedy, UCB, Thompson sampling) are *operational approximations* of AAT primitives under regime-specific scope conditions — none exactly implements the optimal-gain principle, but their gains are *related to* the uncertainty-ratio gain and the methods converge in steady state toward optimal-gain behaviour. The worked example concentrates on the nonstationary case where AAT’s machinery is most distinctive from classical RL: the forgetting prerequisite becomes operational (effective sample size per arm must be bounded or the agent’s gain converges to zero); the per-arm tempo decomposition surfaces persistence as the binding constraint when uniform exploration is mismatched to drift.

System.

A k -armed bandit ($k = 4$) with slowly drifting reward means:

Formulation.

$$\mu_i(t+1) = \mu_i(t) + w_i(t), \quad w_i(t) \sim \mathcal{N}(0, q), \quad q = 0.01 \quad (\text{BA.1})$$

- Observation: Pulling arm $a_t = i$ yields reward $r_t = \mu_i(t) + \varepsilon_t$, $\varepsilon_t \sim \mathcal{N}(0, \sigma^2)$, $\sigma^2 = 1.0$.
- Event rate: $\nu = 1$ pull per step.
- Agent: Q-learning with fixed learning rate α and UCB action selection.

Chain Instantiation.

Scope (1.6). *Mapping: exact.*

$\Omega_t = (\mu_1(t), \dots, \mu_4(t))$ — not directly observable. $\mathcal{A} = \{1, 2, 3, 4\}$. Residual uncertainty persists: means drift continuously, rewards are noisy, and only one arm is observed per step.

Causal Structure (1.7) + CIY (3.7). *Mapping: exact for causal ordering; approximate for CIY formula.*

The action (arm choice) temporally precedes the reward observation. Under the Gaussian reward model, CIY for arm i :

Discussion — Bandit CIY, Gaussian Case.

$$\text{CIY}(i; M_{t-1}) = \frac{1}{2\sigma^2} \mathbb{E}_{j \sim q} [(\hat{\mu}_i - \hat{\mu}_j)^2] \quad (\text{BA.2})$$

This measures how *distinctive* arm i 's predicted reward is relative to alternatives — distinctiveness of predictions, not uncertainty about predictions. A proper uncertainty-aware CIY would require maintaining posterior variances per arm, which the basic Q-learning model does not track.

Model (2.1) and Sufficiency (2.3). *Mapping: exact structural mapping; the model is deliberately impoverished.*

$$M_t = (\hat{\mu}_1, \dots, \hat{\mu}_4, n_1, \dots, n_4) \quad (\text{BA.3})$$

Model sufficiency $S(M_t) < 1$ for two reasons: (1) No drift tracking — the model treats reward means as stationary. (2) No uncertainty representation — unlike a Bayesian agent that would maintain posterior variances.

Mismatch (3.4). *Mapping: exact.*

$$\delta_t = r_t - \hat{\mu}_{a_t} \quad (\text{BA.4})$$

Decomposes exactly: $\mathbb{E}[\delta_t^2] = (\hat{\mu}_{a_t} - \mu_{a_t}(t))^2 + \sigma^2$.

Update Gain (3.6). *Mapping: approximate. This is where the bandit agent is most deficient relative to AAT-optimal behavior.*

The Q-learning update

$$\hat{\mu}_{a_t} \leftarrow \hat{\mu}_{a_t} + \alpha \cdot \delta_t \quad (\text{BA.5})$$

uses a **degenerate constant gain**. It does not adapt to the agent's current uncertainty state.

Optimal gain via effective window. A fixed α is equivalent to exponential discounting with effective window $W = (1 - \alpha)/\alpha$. The per-arm model uncertainty balances estimation variance and drift-induced bias:

$$U_M \approx \frac{\sigma^2}{W} + q \cdot W \quad (\text{BA.6})$$

The optimal window $W^* = \sigma/\sqrt{q} = 10$ gives $\alpha^* \approx 0.091$, $U_M \approx 0.2$, and:

$$\eta^* = \frac{U_M}{U_M + U_o} = \frac{0.2}{1.2} \approx 0.167 \quad (\text{BA.7})$$

The discrepancy between α^* and η^* arises because α operates on raw prediction errors while η^* accounts for the full uncertainty structure. The two converge when the model properly tracks its own uncertainty.

Exploration (3.7). *Mapping: approximate structural analogy.*

UCB: $a_t = \arg \max_i [\hat{\mu}_i + c\sqrt{\ln t / n_i}]$

Table BA.1

<i>AAT component</i>	<i>UCB analog</i>	<i>Quality</i>
$\mathbb{E}[\text{value}(a) \mid M_t]$	$\hat{\mu}_i$	Exact
$\lambda(M_t) \cdot \mathbb{E}[\text{CIY}(a)]$	$c\sqrt{\ln t / n_i}$	Approximate

UCB captures the CIY structure (rarely-visited arms are more informative) but depends on visit count, not observation content.

Tempo + Persistence (3.8, 4.5). *Mapping: approximate.*

Per-arm tempo. With approximately uniform exploration, each arm sampled at rate $\nu/k = 1/4$:

$$\mathcal{T}_i = \frac{\nu}{k} \cdot \alpha = 0.25 \times 0.091 = 0.023 \text{ step}^{-1} \quad (\text{BA.8})$$

Per-arm drift rate: $\rho_i = \sqrt{q} = 0.1 \text{ reward units} \cdot \text{step}^{-1}$.

Persistence check: $\mathcal{T}_i = 0.023$ vs $\rho_i = 0.1$ — **fails**. Per-arm correction tempo is too low relative to drift. The agent cannot keep all arms' estimates current under uniform exploration.

Interpretation. This is expected and informative. AAT diagnoses exactly why: with 4 arms and one pull per step, each arm is visited too infrequently to track its drifting mean. Two remedies:

1. **Increase η^*** (raise α): Extract more per observation, but increase steady-state noise.
2. **Concentrate ν** : Abandon uniform exploration. Focus pulls on a subset, increasing per-arm ν/k_{active} . This is exactly what a good UCB policy does.

With focused exploration ($k_{\text{active}} = 2, \alpha = 0.2$): $\mathcal{T}_i = 0.5 \times 0.2 = 0.1$ — barely meets the threshold. The fundamental tension: with limited pulls and significant drift, the agent must accept either failing to track some arms or using high α that introduces noise.

Aggregate failure. $\mathcal{T} = \nu \cdot \alpha = 0.091$ vs $\rho = k \cdot \sqrt{q} = 0.4$. The agent's total adaptive capacity is outpaced by total environmental drift — a regime where model-based approaches (Bayesian bandits, Kalman bandit filters) with higher effective η^* have a structural advantage.

<i>AAT Concept</i>	<i>Bandit Mapping</i>	<i>Status</i>
--------------------	-----------------------	---------------

Mapping Quality Summary.

Table BA.2

<i>AAT Concept</i>	<i>Bandit Mapping</i>	<i>Status</i>
Scope	Exact	Definitional
Causal structure	Exact	Structural
CIY	Approximate	Model lacks uncertainty representation
Model (M_t)	Exact	Definitional
Model sufficiency ($S < 1$)	Exact	Q-learner demonstrably insufficient
Mismatch (δ_t)	Exact	Standard prediction error
Gain (α as degenerate η^*)	Approximate	Fixed, does not adapt
Exploration (UCB as approximate CIY)	Approximate	Structural analogy
Tempo ($\mathcal{T} = (\nu/k) \cdot \alpha$ per arm)	Approximate	Assumes uniform allocation
Persistence condition	Exact (qualitative)	Threshold structure is robust

Epistemic Status.

This is a *structural mapping*, not a derivation. The mapping is *exact* for scope, causal ordering, model compression, and mismatch. It is *approximate* for gain, exploration, and tempo — precisely because the Q-learner does not represent its own uncertainty, which is what 3.6 requires for optimal behavior. The gap between the Q-learner’s fixed α and the AAT-optimal adaptive η^* is precisely the gap between a fixed-gain PID controller and a Kalman filter.

The mapping status is *conditional* because the quantitative relationships depend on the Gaussian reward model and specific parameterization. The qualitative conclusions (persistence failure under uniform exploration, the concentration-vs-noise tradeoff) should be robust.

Working Notes.

- The per-arm analysis is a natural instance of the per-dimension tempo decomposition (165): each arm is an independent mismatch dimension with its own $\mathcal{T}_i$ and ρ_i . The aggregate tempo overstates effective adaptation along any individual arm’s dimension — exactly the failure mode that the scalar-to-tensor generalization captures.
- A Bayesian bandit (maintaining per-arm posteriors with exponential discounting) would achieve higher 2.3 by representing its own uncertainty, yielding an adaptive η^* that matches the AAT-optimal form. Comparing Q-learner vs Bayesian bandit performance under nonstationarity is a direct test of 3.6’s structural claim.
- Reward-unit and surprise-unit formulations coincide up to normalization by σ^2 . For this example ($\sigma = 1$), they are identical.

Appendix BB

Worked Example: Nonstationary Bandit with Explicit Strategy (Part II Validation)

AAT's Part II machinery — objectives, strategy DAGs, the orient cascade, satisfaction gap, control regret, chain confidence decay, observability dominance, and edge update via gain — is exercised on a concrete system simple enough for closed-form analysis. The key payoff: Part II diagnostics tell the agent things that Part I quantities alone cannot.

System.

A 3-armed Bernoulli bandit with drifting success probabilities and an agent that maintains an explicit strategy DAG for arm selection.

Environment. Three arms with time-varying success probabilities:

Formulation.

$$\theta_k(t) \in [0, 1], \quad k \in \{1, 2, 3\} \quad (\text{BB.1})$$

At each step, pulling arm k yields reward $r_t \in \{0, 1\}$ with $P(r_t = 1 \mid a_t = k) = \theta_k(t)$. The best arm changes slowly: at rate ρ_θ , a random arm's success probability shifts by a random perturbation clipped to $[0.1, 0.9]$.

Concrete parameterization. At $t = 0$: $\theta_1(0) = 0.7$, $\theta_2(0) = 0.5$, $\theta_3(0) = 0.3$. Drift rate: $\rho_\theta = 0.05$ per step (one arm shifts by ± 0.05 every 20 steps on average). Horizon: $N_h = 100$ steps.

Agent. Maintains Beta-Bernoulli beliefs per arm and an explicit strategy DAG encoding its plan for achieving a target reward rate.

Part I Chain Instantiation (Summary).

Part I quantities map as in [BA](#). The brief summary here establishes the epistemic baseline; the novel content is Part II.

Scope (1.6). $\Omega_t = (\theta_1(t), \theta_2(t), \theta_3(t))$; $\mathcal{A} = \{1, 2, 3\}$; residual uncertainty persists. *Exact.*

Model (2.1). $M_t = (\alpha_k, \beta_k)_{k=1}^3$ — Beta posteriors per arm, with exponential discounting (effective window $W = 20$):

$$\hat{\theta}_k = \frac{\alpha_k}{\alpha_k + \beta_k} \quad (\text{BB.2})$$

Exact structural mapping.

Mismatch (3.4). $\delta_t = r_t - \hat{\theta}_{a_t}$. *Exact.*

Update gain (3.6). The Beta-Bernoulli conjugate update with discounted pseudo-counts gives:

$$\eta_k^* = \frac{1}{\alpha_k + \beta_k + 1} \quad (\text{BB.3})$$

For a well-observed arm with $\alpha_k + \beta_k \approx W = 20$: $\eta_k^* \approx 0.048$. *Exact for the pulled arm; zero for unpulled arms.*

Tempo (3.8). Per-arm tempo: $\mathcal{T}_k = \nu_k \cdot \eta_k^*$ where ν_k is the pull rate for arm k . Under uniform exploration: $\nu_k = 1/3$, so $\mathcal{T}_k \approx 0.016$ per step. *Approximate.*

Part II Chain Instantiation.

Objective (6.4). *Mapping: exact.*

The agent's objective: achieve a target average reward rate over the horizon.

Definition (bandit-objective).

$$V_{O_t}(\tau_{0:N_h}) = \frac{1}{N_h} \sum_{t=0}^{N_h-1} r_t \quad (\text{BB.4})$$

$$V_{O_t}^{\min} = 0.6 \quad (\text{BB.5})$$

This is a utility-type objective with a satisfaction threshold: the agent counts the mission as successful if its average reward rate meets or exceeds 0.6. The value functional evaluates trajectories on $[0, 1]$ (average reward).

Value Object (6.6). *Mapping: exact.*

Under the one-step improvement convention ($\pi_{\text{cont}} = \pi_{\text{current}}$):

$$V_O(M_t, \pi; N_h) = \mathbb{E} \left[\frac{1}{N_h} \sum_{s=t}^{t+N_h-1} r_s \mid M_t, \pi \right] \quad (\text{BB.6})$$

For a greedy policy that always pulls the estimated-best arm $k^* = \arg \max_k \hat{\theta}_k$:

$$V_O(M_t, \pi_{\text{greedy}}; N_h) = \hat{\theta}_{k^*} \quad (\text{BB.7})$$

This is tractable: the agent's expected average reward under greedy play equals its belief about the best arm's success probability. With M_0 : $V_O = \hat{\theta}_1 = 0.7$.

Action-value form:

$$Q_O(M_t, k; \pi_{\text{greedy}}, N_h) \approx \frac{1}{N_h} \hat{\theta}_k + \frac{N_h - 1}{N_h} \hat{\theta}_{k^*} \quad (\text{BB.8})$$

The first pull contributes $1/N_h$ of the horizon; the rest follows the greedy policy.

Objective Attainability (8.4). *Mapping: exact.*

$$A_O(M_t; \Pi, N_h) = \sup_{\pi \in \Pi} V_O(M_t, \pi; N_h) \quad (\text{BB.9})$$

For a bandit with Beta-Bernoulli beliefs, the optimal policy (Gittins index or Bayes-adaptive) achieves an attainability that is bounded:

$$\max_k \hat{\theta}_k \leq A_O \leq \max_k \hat{\theta}_k + \epsilon_{\text{info}} \quad (\text{BB.10})$$

where ϵ_{info} accounts for the information value of exploration (pulling a suboptimal arm now to identify the true best arm later). The upper bound is tight when exploration can reveal a better arm. For our initial beliefs: $A_O \approx 0.70 + \epsilon_{\text{info}}$.

Satisfaction gap:

$$\delta_{\text{sat}} = V_{O_t}^{\min} - A_O = 0.6 - 0.70 \approx -0.10 \quad (\text{BB.11})$$

$\delta_{\text{sat}} < 0$: the objective is **attainable**. The agent believes it can achieve a 0.6 average reward rate. The negative gap is margin — the agent has room to spare.

Control Regret (8.5). *Mapping: exact.*

Under an explore-then-exploit policy that spends 30 steps exploring uniformly, then exploits the estimated-best arm:

$$V_O(M_t, \pi_{\text{explore-exploit}}; 100) = \frac{30}{100} \cdot \bar{\theta} + \frac{70}{100} \cdot \hat{\theta}_{k^*(30)} \quad (\text{BB.12})$$

where $\bar{\theta} = (\hat{\theta}_1 + \hat{\theta}_2 + \hat{\theta}_3)/3$ is the average expected reward during exploration and $\hat{\theta}_{k^*(30)}$ is the estimated-best arm after 30 observations. With initial beliefs: $V_O \approx 0.3 \cdot 0.5 + 0.7 \cdot 0.7 = 0.64$.

$$\delta_{\text{regret}} = A_O - V_O(\pi_{\text{explore-exploit}}) \approx 0.70 - 0.64 = 0.06 \quad (\text{BB.13})$$

The agent is leaving 0.06 reward-rate points on the table with its current explore-then-exploit strategy — it explores too much. This is the signal for strategy revision (10.1 step 3).

The 2x2 diagnostic (8.5 Discussion):

Table BB.1

	$\delta_{\text{sat}} \leq 0$	$\delta_{\text{sat}} > 0$
$\delta_{\text{regret}} \approx 0$	Success	Capability limit
$\delta_{\text{regret}} \gg 0$	Strategy problem	Both

Our agent is in the **Strategy problem** cell: the goal is achievable, but the current policy is suboptimal. Part I alone (which reports $\hat{\theta}_k$, η_k^* , and $\mathcal{T}_k$) cannot distinguish “good strategy, hard goal” from “bad strategy, easy goal.” The 2x2 diagnostic is Part II’s value-add.

Strategy DAG (8.3, 8.2). Mapping: exact for structure and propagation; the DAG is simple enough for closed-form status.

The agent's strategy for achieving $V_{O_t}^{\min} = 0.6$:

```
[root: "achieve  $\geq 0.6$  avg reward"] (AND) / \ [Phase 1: "identify best arm"] (OR) [Phase 2: "exploit best arm"]
/ | \ | [arm 1 [arm 2 [arm 3 [pull identified is best] is best] is best] arm reliably]
```

Nodes: - v_{root} : AND-node. Both phases must succeed: identify the best arm AND exploit it. Terminal satisfaction condition: average reward ≥ 0.6 . - v_{id} : OR-node (Phase 1). Identifying the best arm succeeds if ANY one arm is correctly identified as best. - v_1, v_2, v_3 : leaf condition nodes. "Arm k is the best arm." Base credence: $p_k(M_t) = P(\theta_k > \theta_j \forall j \neq k | M_t)$. - v_{exploit} : leaf action node (Phase 2). "Pulling the identified arm yields reward ≥ 0.6 ." Base credence: $p_{\text{exploit}}(M_t) = P(\theta_{k^*} \geq 0.6 | M_t)$.

Edge credences. The edges from leaf nodes to their parents carry causal credences: - $p_{v_k, v_{\text{id}}}$: "correctly identifying arm k as best enables Phase 1 success." These are structural (close to 1.0) — if the agent correctly knows which arm is best, Phase 1 succeeds by definition. Set $p_{v_k, v_{\text{id}}} = 1.0$ for all k . - $p_{v_{\text{id}}, v_{\text{root}}}$: "having identified the best arm advances the root goal." Set to 0.95 (slight uncertainty about whether identification translates to exploitation success). - $p_{v_{\text{exploit}}, v_{\text{root}}}$: "successful exploitation advances the root goal." Set to 0.90 (the identified arm might drift below threshold during exploitation).

Leafbase credences from M_t . At $t = 0$ with $(\hat{\theta}_1, \hat{\theta}_2, \hat{\theta}_3) = (0.7, 0.5, 0.3)$ and moderate pseudo-counts ($\alpha_k + \beta_k \approx 5$) from prior experience:

$$p_1(M_0) = P(\theta_1 > \theta_2, \theta_1 > \theta_3 | M_0) \approx 0.75 \quad (\text{BB.14})$$

$$p_2(M_0) \approx 0.20, \quad p_3(M_0) \approx 0.05 \quad (\text{BB.15})$$

$$p_{\text{exploit}}(M_0) = P(\theta_{k^*} \geq 0.6 | M_0) \approx 0.70 \quad (\text{BB.16})$$

Status propagation (8.3, 8.2):

Phase 1 (OR-node):

$$s_{v_{\text{id}}} = 1 - \prod_{k=1}^3 (1 - p_{v_k, v_{\text{id}}} \cdot p_k) = 1 - (1 - 0.75)(1 - 0.20)(1 - 0.05) \quad (\text{BB.17})$$

$$= 1 - (0.25)(0.80)(0.95) = 1 - 0.19 = 0.81 \quad (\text{BB.18})$$

Phase 2 (leaf): $s_{v_{\text{exploit}}} = 0.70$.

Root (AND-node):

$$s_{v_{\text{root}}} = (p_{v_{\text{id}}, v_{\text{root}}} \cdot s_{v_{\text{id}}}) \cdot (p_{v_{\text{exploit}}, v_{\text{root}}} \cdot s_{v_{\text{exploit}}}) \quad (\text{BB.19})$$

$$= (0.95 \times 0.81) \times (0.90 \times 0.70) = 0.770 \times 0.630 = 0.485 \quad (\text{BB.20})$$

Plan-confidence score: $\hat{P}_2(M_0) = 0.485$. The strategy assigns about 49% confidence to achieving the objective. This is distinct from A_O (which optimizes over all policies) — $\hat{P}_2$ is the DAG's self-assessment of *this specific plan*.

Chain Confidence Decay (8.1). *Mapping: exact (mathematical identity).*

The root node's confidence illustrates chain decay through an AND structure. The root requires two children to succeed, each with their own sub-chains:

$$\log \hat{R}_2 = \log(p_{v_{\text{id}}, u_{\text{root}}} \cdot s_{v_{\text{id}}}) + \log(p_{v_{\text{exploit}}, u_{\text{root}}} \cdot s_{v_{\text{exploit}}}) \quad (\text{BB.21})$$

$$= \log(0.770) + \log(0.630) = -0.261 + (-0.462) = -0.724 \quad (\text{BB.22})$$

Each term is negative; the AND combination accumulates them additively in log-space. If the agent added a third AND-child (e.g., “maintain budget”), even with $s_{\text{budget}} = 0.95$, the strategy-plan confidence would drop to $0.485 \times 0.95 = 0.461$. The pressure toward shallow, narrow DAGs is structural.

Contrast with the OR-node. Phase 1 (OR-node) benefits from redundancy: three alternative arms provide resilience. The OR combination yields $s_{v_{\text{id}}} = 0.81$ even though no single arm has credence above 0.75. The asymmetry between AND (multiplicative decay) and OR (complementary resilience) is the core structural result of 8.1 combined with 8.2.

Orient Cascade (10.1). *Mapping: exact for ordering; the content of each step is tractable in this domain.*

Suppose at $t = 40$, the agent has been pulling arm 1 (the initially-best arm) and observes a run of failures: the last 8 of 10 pulls yielded $r_t = 0$. Meanwhile, arm 2's true probability has drifted to $\theta_2(40) = 0.8$ (unknownst to the agent) and arm 1 has drifted to $\theta_1(40) = 0.3$.

The orient cascade proceeds:

Step 1: Reduce $\delta_{\text{epistemic}}$. Update M_t from the recent observations. The Beta posterior for arm 1 shifts:

$$\hat{\theta}_1 : 0.70 \rightarrow 0.38 \quad (\text{after 8 failures in 10 trials with discounted prior}) \quad (\text{BB.23})$$

Arms 2 and 3 are unobserved since the exploration phase ended. Their posteriors remain at approximately $\hat{\theta}_2 = 0.52$, $\hat{\theta}_3 = 0.28$ (slightly regressed toward the prior due to discounting). *The epistemic update proceeds without reference to O_t or Σ_t* — directed separation (6.3) holds because the Beta update depends only on the observation and the prior, not on the agent's goals.

Step 2: Evaluate δ_{sat} . With the updated M_t :

$$A_O(M_{40}) \approx \max_k \hat{\theta}_k + \epsilon_{\text{info}} = 0.52 + \epsilon_{\text{info}} \quad (\text{BB.24})$$

$$\delta_{\text{sat}} = 0.6 - (0.52 + \epsilon_{\text{info}}) \quad (\text{BB.25})$$

If $\epsilon_{\text{info}} \approx 0.05$ (exploration can reveal a better arm): $\delta_{\text{sat}} \approx 0.6 - 0.57 = 0.03 > 0$. The objective is now **marginally unmet** — the agent believes it may not be able to achieve 0.6 average reward. This triggers the disambiguation procedure (8.4): is the goal truly infeasible, or is M_t wrong about what's achievable?

Step 3: Evaluate δ_{regret} . Under the current greedy policy (still pulling arm 1 despite updated beliefs):

$$V_O(M_{40}, \pi_{\text{current}}) = \hat{\theta}_1 = 0.38 \quad (\text{BB.26})$$

$$\delta_{\text{regret}} = 0.57 - 0.38 = 0.19 \quad (\text{BB.27})$$

Control regret is high. The agent is pulling arm 1 out of inertia, but M_t now says arm 2 is better. *This is the Part II payoff*: Part I would report the mismatch δ_t on each pull and the updated $\hat{\theta}_1$, but it cannot diagnose that the *strategy* is the problem. The high δ_{regret} identifies the corrective action: revise Σ_t , not O_t .

Step 4: Evaluate $\delta_{\text{strategic}}$. The edge residual for the root AND-node's exploitation branch:

$$r_{\text{exploit}} = \mathbb{E}[\Delta V_O \mid \text{exploit arm 1}] - \Delta V_O^{\text{observed}} \quad (\text{BB.28})$$

The agent predicted reward ≈ 0.70 per pull from arm 1; it observed ≈ 0.20 . The residual is large and positive, localizing the problem to the exploitation edge: the agent's causal belief that "pulling arm 1 yields high reward" is wrong.

Simultaneously, the leaf credences update: $p_1(M_{40}) \approx 0.15$ (arm 1 is probably not the best), $p_2(M_{40}) \approx 0.60$ (arm 2 is probably best, by default from the prior). The Phase 1 OR-node status shifts: the identification sub-goal may need to be revisited.

Step 5: Revise Σ_t . The agent restructures its strategy: switch exploitation target to arm 2 (the currently estimated-best arm). In DAG terms: $p_1(M_{40})$ has dropped to 0.15 while $p_2(M_{40})$ has risen to 0.60. The OR-node now routes through arm 2. The exploitation leaf updates: $p_{\text{exploit}}(M_{40}) = P(\theta_2 \geq 0.6 \mid M_{40}) \approx 0.40$.

After strategy revision: $\hat{P}_2(M_{40}) \approx (0.95 \times 0.73) \times (0.90 \times 0.40) = 0.693 \times 0.360 = 0.249$. Plan confidence has dropped from 0.485 to 0.249 — the agent correctly recognizes that its situation has deteriorated.

What Part I alone cannot provide. At $t = 40$, Part I reports: $\hat{\theta}_1 = 0.38$, $\hat{\theta}_2 = 0.52$, $\hat{\theta}_3 = 0.28$, and the per-arm mismatch signals. This tells the agent what it believes about each arm. It does NOT tell the agent: - Whether its *goal* is still achievable (δ_{sat}) - Whether its *strategy* is the problem or the goal is (δ_{regret}) - Which *part* of the strategy is failing ($\delta_{\text{strategic}}$, edge residuals) - What corrective action to take (the cascade's ordering)

Observability Dominance (9.3). *Mapping: exact mechanism; approximate quantitative form.*

When the agent stops pulling arm 3 after the exploration phase (say at $t = 30$), the observation rate for arm 3 drops to $\nu_3 = 0$. By 3.6:

$$\eta_3^* = \frac{U_{\text{edge},3}}{U_{\text{edge},3} + U_{\text{obs},3}} \quad (\text{BB.29})$$

With $U_{\text{obs},3} \rightarrow \infty$ (no observations): $\eta_3^* \rightarrow 0$. Arm 3's credence $p_3(M_t)$ is **frozen at its exploration-phase value**.

For the strategy DAG, the leaf node "arm 3 is best" has: - Nominal credence: $p_3(M_t) \approx 0.05$ (from $t = 30$ beliefs) - Path observability: $\sigma_3 = 0$ (arm 3 is not being pulled) - Observability-adjusted confidence: $\text{conf}_{\text{obs}}(\text{arm 3 path}) = 0.05 \times 0 = 0$

Arm 3's path through the strategy is **epistemically dead** — the agent cannot learn whether arm 3 has become the best arm. If θ_3 drifts to 0.9 (which it might, given nonstationarity), the agent will never discover this through the exploitation-only policy.

The observability-dominance diagnostic. Compare arm 2 and arm 3:

Table BB.2

	Arm 2	Arm 3
Nominal credence p_k	0.52	0.28 (frozen)
Observability σ_k	moderate (occasional pulls)	0 (never pulled)
Adjusted confidence	$0.52 \times \sigma_2$	0
Can be updated?	Yes	No

Arm 2 is epistemically alive; arm 3 is epistemically dead. The agent should prefer exploring arm 2 (where it can learn) over arm 3 — unless it invests in making arm 3 observable again (pulling it despite low expected reward). This is the strategic case for exploration: not just expected value, but *observability maintenance*.

Part I vs Part II diagnosis. Part I reports $\hat{\theta}_3 = 0.28$ with growing uncertainty (the posterior variance increases under discounting with no new data). Part II adds the diagnostic: this arm's strategy path is epistemically dead ($\sigma_3 = 0$), and restoring observability has value beyond the arm's nominal expected reward.

Edge Update via Gain (9.4). *Mapping: exact for the Beta-Bernoulli case.*

When the agent pulls arm 1 at $t = 35$ and observes failure ($r_t = 0$):

Leaf credence update. The Beta posterior for arm 1 updates:

$$\beta_1 \rightarrow \beta_1 + 1, \quad \hat{\theta}_1 = \frac{\alpha_1}{\alpha_1 + \beta_1 + 1} \quad (\text{BB.30})$$

This is the standard conjugate update — the gain principle reduces to Bayesian updating in the binary case (9.4).

Strategy propagation. The updated $\hat{\theta}_1$ changes $p_1(M_t)$ (the credence that arm 1 is best). This propagates through the OR-node to the root:

$$s_{v_{\text{id}}}^{\text{new}} = 1 - (1 - p_1^{\text{new}})(1 - p_2)(1 - p_3) \quad (\text{BB.31})$$

The OR-node's status is *resilient* to single-arm credence drops: even if p_1 drops substantially, p_2 can compensate. This is the OR-node's structural advantage — redundancy absorbs single-edge failures.

The exploitation edge. The edge $p_{v_{\text{exploit}}, v_{\text{root}}} = 0.90$ is a causal credence: “pulling the identified-best arm reliably yields reward ≥ 0.6 .” After observing repeated failures from arm 1:

$$\text{signal} = 0 \quad (\text{failure observed}) \quad (\text{BB.32})$$

$$p_{\text{exploit}}^{\text{new}} = p_{\text{exploit}}^{\text{old}} + \eta_{\text{edge}} \cdot (0 - p_{\text{exploit}}^{\text{old}}) \quad (\text{BB.33})$$

With $\eta_{\text{edge}} = U_{\text{edge}} / (U_{\text{edge}} + U_{\text{obs}})$, the exploitation edge's credence decreases. The rate of decrease is governed by the gain principle: well-established edges (many prior successes, U_{edge} small) are slow to revise; newly formed edges are fast.

Directed Separation (6.3). *Mapping: exact. The agent is Class 1 (Separated) by construction.*

The Beta-Bernoulli update for M_t processes (α_t, r_t) without reference to (O_t, Σ_t) :

$$(\alpha_{a_t}, \beta_{a_t}) \rightarrow (\alpha_{a_t} + r_t, \beta_{a_t} + 1 - r_t) \quad (\text{BB.34})$$

The update function f_M has no G_t argument. The strategy DAG then updates its leaf credences from the new M_t — this is $f_G(G_t, M_{t+1}, e_t)$. The agent architecture is modular: a separate estimator (Beta posteriors) and planner (DAG evaluation), connected through the state-estimate interface.

$\kappa_{\text{processing}} = 0$ — no goal information flows into the epistemic update. Part II results apply exactly.

Satisfaction Gap Dynamics. *Mapping: exact definition; approximate computation.*

Tracking δ_{sat} over time reveals the agent’s evolving assessment of goal feasibility:

Table BB.3

Time	$\hat{\theta}_{k^*}$	A_O (approx)	δ_{sat}	Interpretation
$t = 0$	0.70	0.72	−0.12	Comfortably attainable
$t = 20$	0.68	0.70	−0.10	Still attainable
$t = 40$	0.52	0.57	+0.03	Marginally unmet
$t = 50$ (after exploring arm 2)	0.75	0.77	−0.17	Attainable again — M_t improved
$t = 70$	0.78	0.79	−0.19	Confidently attainable

The key transition is $t = 40 \rightarrow 50$: δ_{sat} goes from positive (unmet) to negative (attainable) — not because the goal changed or the environment improved, but because M_t improved. The agent explored arm 2, discovered $\hat{\theta}_2 \approx 0.75$, and its attainability assessment corrected upward. This matches the 8.4 disambiguation: the positive δ_{sat} at $t = 40$ was caused by a bad model, not an infeasible goal. The orient cascade’s ordering (fix M_t first, then reassess δ_{sat}) produced the right diagnosis.

The Part II Payoff.

The purpose of this worked example is to demonstrate that Part II adds genuine diagnostic value beyond Part I. Here is the explicit comparison:

Scenario at $t = 40$: arm 1 failing, arm 2 actually best, agent pulling arm 1.

Part I diagnosis. The mismatch signal δ_t is large. The posterior $\hat{\theta}_1$ has dropped. Per-arm tempo $\mathcal{T}_1 \approx 0.016$ is adequate for tracking arm 1’s drift. The agent knows arm 1 looks bad. Part I prescribes: update M_t faster (increase η^*) or observe more (increase ν).

Part II diagnosis. The orient cascade runs: 1. M_t updates correctly — arm 1’s posterior drops to 0.38. (Directed separation: this happens goal-blindly.) 2. $\delta_{\text{sat}} = +0.03$ — the objective is marginally unmet. The agent’s best-believed policy barely misses the target. 3. $\delta_{\text{regret}} = 0.19$ — the current policy (pulling arm 1) is far from the best available (pulling arm 2). 4. The 2x2 diagnostic says: “goal may be achievable, strategy is definitely suboptimal.” Corrective action: revise Σ_t , then reassess δ_{sat} . 5. Edge residuals localize the problem: the exploitation branch’s predicted-vs-observed value gap is large. The leaf credences identify arm 2 as the best candidate. 6. Observability check: arm 3 is epistemically dead ($\sigma_3 = 0$). The agent should invest in observing arm 2 (or arm 3) before committing.

Part II tells the agent: *switch to arm 2, but first explore to confirm — the goal is probably still achievable, the current strategy is the problem, not the goal*. Part I tells the agent: *arm 1 looks bad*. The difference is the difference between a diagnosis and a symptom.

Mapping Quality Summary.

Table BB.4

<i>AAT Concept</i>	<i>Bandit-Strategy Mapping</i>	<i>Status</i>	<i>Notes</i>
Scope (1.6)	Exact	Definitional	Same as BA
Model M_t (2.1)	Exact	Structural	Beta posteriors per arm
Mismatch δ_t (3.4)	Exact	Standard	Prediction error
Update gain η^* (3.6)	Exact	Conjugate	Beta-Bernoulli gain
Tempo $\mathcal{T}$ (3.8)	Approximate	Per-arm decomposition	Assumes known allocation
Directed separation (6.3)	Exact	Class 1 (Separated) by construction	Beta update is goal-blind
Objective O_t (6.4)	Exact	Utility with threshold	$V_{O_t}^{\min} = 0.6$
Value object V_O (6.6)	Exact	Closed-form (greedy)	$V_O = \hat{\theta}_{k^*}$
Satisfaction gap δ_{sat} (8.4)	Exact	Definition applied	$V_{O_t}^{\min} - A_O$
Control regret δ_{regret} (8.5)	Exact	Definition applied	$A_O - V_O(\pi_{\text{current}})$
Strategy DAG Σ_t (8.3)	Exact	AND/OR with 7 nodes	Small enough for manual propagation
AND/OR scope (8.2)	Exact	Natural fit	Phase 1 is genuinely OR; root is genuinely AND
Status propagation	Exact	Closed-form	Forward pass computable by hand
Plan confidence $\hat{P}_\Sigma$ (8.3)	Exact (computation)	From propagation; systematically optimistic under correlated failure	0.485 at $t = 0$
Chain confidence decay (8.1)	Exact	Mathematical identity	AND-node multiplicative; OR-node resilient
Orient cascade (10.1)	Exact (ordering)	All 5 steps exercised	Content is tractable in this domain
Observability dominance (9.3)	Exact (mechanism)	Unpulled arms frozen	$\sigma_3 = 0 \Rightarrow \eta_3^* = 0$
Edge update via gain (9.4)	Exact (binary case)	Beta-Bernoulli conjugate	Gain principle = Bayesian update
Strategic calibration (9.1)	Approximate	Credit assignment tractable here	Single-arm attribution is clean
Satisfaction gap dynamics	Exact	Time series of δ_{sat}	Model improvement resolves positive gap

Quantities that map cleanly. All Part II definitions (objective, value object, satisfaction gap, control regret, strategy DAG) have exact instantiations. Plan confidence $\hat{P}_\Sigma$ is exact as a computation (the AND/OR propagation is correct), but systematically overestimates actual success probability under correlated edge failures (8.3). In this toy bandit, edge independence is a reasonable approximation; in complex domains the overestimate may be severe. The orient cascade ordering is exact and all five steps are exercised.

Quantities that map approximately. Strategic calibration is approximate because even in this simple domain, attributing value changes to specific DAG edges requires assumptions about the agent's execution fidelity. Observability dominance is exact in mechanism but the multiplicative form $\text{conf}_{\text{obs}} = \text{conf} \cdot \text{obs}$ is a first-order approximation.

Quantities that don't map. No Part II quantity fails to map in this domain. The system is simple enough that every quantity is computable (at least approximately). This is the advantage of the bandit setting: it is the simplest system that exercises all Part II machinery.

Epistemic Status.

This is a *worked instantiation* — it demonstrates that the Part II formal chain is internally consistent and instantiable. The mapping quality is *conditional*: the quantitative relationships depend on the Beta-Bernoulli reward model, the specific DAG structure, and the parameterization. The qualitative conclusions — especially the diagnostic value of the 2x2 table and the orient cascade's ordering — should be robust across reward models and DAG topologies.

The status is *conditional* rather than *exact* because the attainability A_O is approximated (the exact Bayes-optimal policy for a nonstationary bandit is intractable), and the strategic calibration mapping inherits the discussion-grade status of 9.1. The core Part II definitions (satisfaction gap, control regret, strategy-plan confidence) are exact instantiations.

Working Notes.

- This example uses a deliberately simple DAG (7 nodes, depth 2). A richer example with depth 3+ would better exercise chain confidence decay's exponential penalty and the evidence-starvation effect from 9.3's quantitative discussion. A natural extension: add an intermediate "confirm arm k is best" node between identification and exploitation, creating a 3-level chain.
 - The explore-then-exploit policy is deliberately suboptimal to demonstrate nonzero δ_{regret} . A Gittins-index agent would achieve $\delta_{\text{regret}} \approx 0$, demonstrating the "Success" cell of the 2x2 table. Comparing the two in the same environment would validate the control-regret diagnostic.
 - The satisfaction gap dynamics table (Section "Satisfaction Gap Dynamics") would benefit from simulation confirmation. The specific numbers are plausible but hand-computed. A Monte Carlo study averaging over environment drift realizations would strengthen the claim.
 - This example does not exercise strategy-level tempo (9.10) or the three-way tradeoff (10.2). A future worked example targeting these would need a system with genuine deliberation costs — simulation in `spikes/sim-three-way-tradeoff.py` provides a drifting-bandit version but found deliberation rarely dominates in simple settings.
 - The DAG's AND/OR assignment is unambiguous in this domain (Phase 1 is genuinely disjunctive; the root is genuinely conjunctive). Domains where the AND/OR assignment is uncertain or updatable would be a stronger test of 8.2.
-

Appendix BC

Worked Example: L1 Augmented DAG with Common-Cause Node

A concrete instantiation of the Correlation Hierarchy (8.3) and Proposition B.6 (R). Two OR-alternatives sharing an infrastructure dependency are modeled at L0 (independence) and L1 (augmented DAG), with the sector condition verified for L1.

The example **exhibits AAT’s central observation about causally-insufficient DAGs**. At L0 (independence), AND/OR propagation systematically *overestimates* plan success: the strategy-plan-confidence comes out higher than the actual rate because the model treats the two paths as independently failing rather than as both depending on shared infrastructure. The overestimate is *exactly* the covariance between the paths’ outcomes — the direction-of-bias formula for OR-structures. At L1 (augmented DAG with the infrastructure node factored *above* the correlated OR-alternatives as a strict-prerequisite AND-node), propagation is correct: conditional on infrastructure up, the paths are independent and standard propagation gives the exact answer. The L1 construction principle (factor the common cause above the correlation it creates) recovers the correct probability. The sector parameter under L1 augmentation has a three-way gating form — condition-testing rate, evidence-starvation on each path under θ_C , exploration-gating on the OR-node — with each factor reflecting a different mechanism. When the common cause is *unobservable*, the L1 augmentation breaks and 5.1 Instance 2 kicks in: the per-conditional decomposition is Cramér-Rao-floored. The agent’s three escape routes (augment observability; multi-child joint observation; fall back to L0-on-marginals) are exercised below with their respective trade-offs.

Setup.

Two alternative paths to a goal G , sharing a common infrastructure dependency C :

- Infrastructure up with probability $\theta_C = 0.8$
- Path 1 succeeds with probability $\theta_{1|C} = 0.9$ when infrastructure is up, fails when down
- Path 2 succeeds with probability $\theta_{2|C} = 0.7$ when infrastructure is up, fails when down
- Goal succeeds if at least one path succeeds (OR)

Actual plan success probability (from the generative model):

$$P(G) = \theta_C \cdot [1 - (1 - \theta_{1|C})(1 - \theta_{2|C})] = 0.8 \cdot 0.97 = 0.776 \quad (\text{BC.1})$$

L0 Analysis: The Independence Model.

The L0 DAG has two leaf nodes with marginal success probabilities:

$$\theta_1 = \theta_C \cdot \theta_{1|C} = 0.72, \quad \theta_2 = \theta_C \cdot \theta_{2|C} = 0.56 \quad (\text{BC.2})$$

OR-propagation:

$$\hat{P}_2^{L0} = 1 - (1 - \theta_1)(1 - \theta_2) = 1 - (0.28)(0.44) = 0.877 \quad (\text{BC.3})$$

Overestimation: $0.877 - 0.776 = 0.101$ (13% relative error).

The overestimation equals $\text{Cov}(X_1, X_2)$ — the covariance induced by the shared infrastructure. The OR-node’s redundancy is illusory: both paths fail together when infrastructure fails, so the “backup” provides less protection than the independence model predicts. This is the $-\rho$ bias from the Correlation Hierarchy’s direction-of-bias formula.

Naive L1: Why Adding a Node Is Not Sufficient.

The obvious L1 construction adds C as a parent of both paths:

- $B_1 = \text{AND}(C, A_1), B_2 = \text{AND}(C, A_2)$
- $G = \text{OR}(B_1, B_2)$

Status propagation: $s_{B_1} = 0.8 \cdot 0.9 = 0.72, s_{B_2} = 0.8 \cdot 0.7 = 0.56, s_G = 1 - (0.28)(0.44) = 0.877$.

Same answer as L0. The OR-propagation treats B_1 and B_2 as marginally independent, but they are correlated through their shared parent C . Adding the common-cause node does not fix the overestimation because the node is *inside* the correlated structure (below the OR-node), not factored above it.

Correct L1: Factor the Common Cause Above the Correlation.

The correct L1 restructuring places C as an AND-prerequisite *above* the OR-structure:

- $G_{\text{sub}} = \text{OR}(A_1, A_2)$ — the sub-plan assuming infrastructure holds
- $G = \text{AND}(C, G_{\text{sub}})$ — infrastructure must hold AND sub-plan must work

Status propagation:

$$s_{G_{\text{sub}}} = 1 - (1 - \theta_{1|C})(1 - \theta_{2|C}) = 1 - (0.1)(0.3) = 0.97 \quad (\text{BC.4})$$

$$s_G = \theta_C \cdot s_{G_{\text{sub}}} = 0.8 \cdot 0.97 = 0.776 \quad (\text{BC.5})$$

Correct. Matches the actual plan success probability exactly.

Why this works: C is above the OR-node, so its value is factored out before the OR-propagation occurs. Conditional on $C = 1$, the two paths A_1 and A_2 are genuinely independent (their exogenous noise terms are independent — the common cause has been conditioned away). The OR-propagation’s independence assumption holds within G_{sub} because the correlation source is above, not among the siblings.

The L1 construction principle: *factor the common cause above the correlation it creates.* When a common cause correlates siblings under an OR-node, restructure so the common cause gates the entire OR-structure via an AND-relationship. This ensures conditional independence of the OR-children and makes standard propagation correct.

Sector Condition Verification (Proposition B.6).

Three leaf nodes with Beta-Bernoulli dynamics and ε -greedy path selection:

Table BC.1

<i>Leaf</i>	<i>True value</i>	<i>Tested when</i>	<i>Expected correction</i>
C (condition)	$\theta_C = 0.8$	Every trial	$-\delta_C/(n_C + 1)$
A_1 (greedy)	$\theta_{1 C} = 0.9$	$C = 1$ AND path 1 selected	$-\theta_C(1 - \varepsilon) \delta_{A_1}/(n_{A_1} + 1)$
A_2 (explore)	$\theta_{2 C} = 0.7$	$C = 1$ AND path 2 selected	$-\theta_C \varepsilon \delta_{A_2}/(n_{A_2} + 1)$

The sector parameter (full proof in Prop B.6 of [R](#)):

$$\alpha_\Sigma = \min\left(\frac{1}{n_C + 1}, \frac{\theta_C(1 - \varepsilon)}{n_{A_1} + 1}, \frac{\theta_C \varepsilon}{n_{A_2} + 1}\right) \quad (\text{BC.6})$$

Three-way gating. This is the first verified mixed AND/OR topology, combining:

1. **Condition testing** — C is tested every trial, same rate as a single edge (B.1)
2. **Evidence starvation** — action edges are gated by θ_C (tested only when $C = 1$), same mechanism as B.2
3. **Exploration gating** — OR-alternatives compete for test opportunities (ε vs $1 - \varepsilon$), same mechanism as B.4

The B.5b bridge applies (non-negative Jacobian, componentwise corrections): $\alpha_s = \alpha_\Sigma$ with no penalty.

L0 vs L1 Comparison.

Table BC.2

<i>Quantity</i>	<i>L0</i>	<i>L1 (correct)</i>
Leaf edges	2 (A_1, A_2 at marginals)	3 ($C, A_1 \mid C, A_2 \mid C$)
Φ at truth	0.877 (biased)	0.776 (correct)
Φ – actual success	+0.101 (overestimates)	0 (unbiased)
Bottleneck α_Σ (at $\varepsilon = 0.3$)	$0.3/(n_2 + 1)$	$0.24/(n_{A_2} + 1)$
Persistence threshold	Lower (easier)	Higher (harder)

The tradeoff: L0 has a higher sector parameter (easier to maintain persistence) but its reference value Φ^{L0} is systematically wrong — it overestimates actual success by the covariance ρ . L1 has a lower sector parameter (the θ_C attenuation makes action calibration harder) but its reference value Φ^{L1} is correct.

Both models have $\delta_s = 0$ when credences are at truth. The difference is what $\delta_s = 0$ means:

- **L0**: well-calibrated within the independence model — which overstates plan success by ρ
- **L1**: well-calibrated to actual plan success

L1 persistence is harder but calibration is honest. The extra leaf (C) adds a degree of freedom, and the conditioning slows action calibration. But the resulting target (Φ^{L1}) is the true plan success probability, not a biased surrogate.

When Correct L1 Construction Is Not Possible.

The factoring-above principle works under two conjoint conditions: (a) the common cause cleanly gates all correlated children through a single AND-relationship (*topological*), and (b) the common cause is a *strict prerequisite* — $\theta_{\text{child}|\neg C} \approx 0$, so children cannot succeed when the common cause is absent (*semantic*). Both conditions must hold. The example above satisfies both: infrastructure gates all alternatives, and absent infrastructure forces failure ($\theta_{i|\neg C} = 0$).

Topological failure: mixed gating. When some OR-alternatives share a common cause and others do not, the factoring-above construction cannot place C above the full sub-plan. Conditioning-based propagation:

$$P(G) = \sum_c P(C = c) \cdot R_\Sigma(G \mid C = c) \quad (\text{BC.7})$$

where $R_\Sigma(G \mid C = c)$ is computed by standard AND/OR propagation with C fixed. Cost: $O(2^k \cdot (|V| + |E|))$ where k is the number of common-cause nodes. Tractable for small k ; exponential in general.

Semantic failure: soft facilitators. When the common cause is a soft facilitator — $\theta_{\text{child}|\neg C} > 0$, so children can still succeed when C is absent, just less reliably (favorable market conditions, a supportive environment, an enabling but non-essential resource) — the AND-prerequisite construction mathematically forces $P(\text{sub-plan} \mid \neg C) = 0$, which strictly understates actual success probability. This is a real failure mode even when the topology admits factoring-above. The natural repair at $O(1)$ cost per soft-facilitator node is the **mixture form (L1')**:

$$\hat{R}_\Sigma^{L1'} = \theta_C \cdot R_\Sigma(G \mid C) + (1 - \theta_C) \cdot R_\Sigma(G \mid \neg C) \quad (\text{BC.8})$$

where the two conditional sub-plans may share structure but carry separate edge credences $p_{ij|C}$ and $p_{ij|\neg C}$. The parametric cost per soft-facilitator-affected edge doubles; the propagation remains polynomial. This is the gap between L1 (strict prerequisites) and L2 (arbitrary joint): L1' handles soft facilitators without the $O(2^k)$ exponential blowup of L2's explicit conditioning. Sector-condition verification for L1' is derived as **Proposition B.7** in *R when C is observable per trial*: under Beta-Bernoulli updating, componentwise updates per conditional branch, and facilitator monotonicity ($R_\Sigma(G|_C) \geq R_\Sigma(G|_{\neg C})$), the sector condition holds globally with five-way gating

$$\alpha_{L1'} = \min\left(\frac{1}{n_C+1}, \min_{j \in \mathcal{J}_C} \frac{\theta_C \pi_{j|C}}{n_{j|C}+1}, \min_{j \in \mathcal{J}_{\neg C}} \frac{(1-\theta_C) \pi_{j|\neg C}}{n_{j|\neg C}+1}\right) \quad (\text{BC.9})$$

(condition testing; evidence starvation and exploration gating on each of the $C = 1$ and $C = 0$ branches). The B.5b transfer to plan value is lossless by facilitator monotonicity plus AND/OR monotonicity on each conditional sub-DAG. In the strict-prerequisite limit $\theta_{j|\neg C} \rightarrow 0$, B.7 reduces to B.6's three-way gating. **When C is not observable on a single observation channel**, the mixture is refuted as a no-go: the per-trial Fisher information matrix on $(\theta_C, p_{j|C}, p_{j|\neg C})$ factors as rank 1 around the indeterminacy manifold $\{\hat{\theta}_C \hat{p}_{j|C} + (1 - \hat{\theta}_C) \hat{p}_{j|\neg C} = \mu_j\}$, so by the Cramér-Rao bound no unbiased online estimator on the joint conditional vector admits $\alpha > 0$. Repair routes when C is unobservable: (i) augment C -observability (recovers B.7); (ii) run $K \geq 2$ children jointly under the same C -realization so the joint Fisher matrix reaches rank $2K + 1$; or (iii) fall back to plan-level tracking on the marginal $\hat{\mu}_j$ (recovers B.1 at the scalar marginal, losing the per-conditional decomposition). See 5.1 for the meta-pattern shared with the on-policy detection no-go in 9.2.

For strategies with many mixed or soft common causes, the agent faces a practical choice: model the most important common causes at L1 (for strict prerequisites) or L1' (for soft facilitators), accept residual L0 bias from unmodeled causes, or invest in richer propagation algorithms. See 8.3 Correlation Hierarchy for the higher-level treatment of L1 / L1' / L2.

Detecting Latent Common Causes.

The detection of causal insufficiency and interventional localization of latent common causes is treated as a standalone result in 9.2 — a *no-go theorem*: under purely on-policy short-circuit execution, no statistic of the agent's observable history can detect the latent, and the L0 residual in particular is *identically zero* on-policy. The residual figures of this example are the **marginal-sampling** regime ($\epsilon = 1$: every trial, both paths observed regardless of policy): there the credences converge to the marginals θ_1, θ_2 and the L0 residual $\Phi^{L0} - \bar{y}_G$ converges to $+\rho$ (our example: $0.877 - 0.776 = 0.101$). Between the two regimes the residual scales as $O(\epsilon)$ and serves only as a *confirmatory* signal; the primary detector is the pairwise sibling covariance test under exploration (9.2).

On-policy observational equivalence (the no-go's shallow-case verification). This example is also the shallow case for which the no-go is *exact*, and the matched-world construction can be verified directly. The on-policy world $\mathcal{W}_{L0}^*$ has independent edges at the regime conditionals

$$\theta_1^* = \theta_C \theta_{1|C} = 0.72, \quad \theta_2^* = \frac{\theta_C (1 - \theta_{1|C}) \theta_{2|C}}{1 - \theta_C \theta_{1|C}} = \frac{0.056}{0.28} = 0.2. \quad (\text{BC.10})$$

Under sequential short-circuit OR execution (try path 1; try path 2 only on failure), the three on-policy observable events have identical probabilities in both worlds:

Table BC.3

Event	$\mathcal{W}_{L1}$ (latent C)	$\mathcal{W}_{L0}^*$ (independent)
A_1 succeeds	0.72	$\theta_1^* = 0.72$
A_1 fails, A_2 succeeds	$0.8 \cdot 0.1 \cdot 0.7 = 0.056$	$0.28 \cdot 0.2 = 0.056$
A_1 fails, A_2 fails	$0.2 + 0.8 \cdot 0.1 \cdot 0.3 = 0.224$	$0.28 \cdot 0.8 = 0.224$

The agent's on-policy history is drawn from the same distribution in both worlds — no on-policy statistic can tell them apart, which is the no-go's content instantiated on this example's numbers.

Epistemic Status.

Conditional on Beta-Bernoulli model. The L0 overestimation, L1 correction, and sector condition verification are exact algebra. The direction-of-bias formula ($\pm\rho$) is exact for two binary siblings **under marginal sampling** (credences at the marginals θ_j ; under on-policy short-circuit execution the credences converge to regime conditionals instead and the residual vanishes — see 9.2); for deeper DAGs with multiple common causes, the bias involves higher-order covariance terms. The on-policy observational-equivalence table is exact arithmetic on the matched-world construction of 9.2. The L1 construction principle (factor common cause above the correlation) is general for single common causes gating OR-siblings; the conditioning-based extension for more complex topologies is standard (variable elimination in Bayesian networks).

Appendix BD

Worked Example: #worked-example-cam (missing)

Segment worked-example-cam not yet written.

Claim: Coevolving automata (Miller 2022): AAT $\leftrightarrow$ Moore machine mapping, meta-machine as $\varepsilon^* = 0$ composition, simplest adaptive agent

Bibliography

- [1] Elias Bareinboim et al. “On Pearl’s hierarchy and the foundations of causal inference”. In: *Probabilistic and Causal Inference: The Works of Judea Pearl*. Ed. by Hector Geffner, Rina Dechter, and Joseph Y. Halpern. ACM Books, 2022, pp. 507–556. DOI: [10.1145/3501714.3501743](https://doi.org/10.1145/3501714.3501743) (cit. on pp. 263, 267, 272, 276, 294, 296).
- [2] Judea Pearl. *Causality: Models, Reasoning, and Inference*. 2nd. Cambridge University Press, 2009 (cit. on pp. 267, 294–296, 298).
- [3] J. Pearl. “A Probabilistic Calculus of Actions”. In: abs/1302.6835 (1994). DOI: [10.1016/B978-1-55860-332-5.50062-6](https://doi.org/10.1016/B978-1-55860-332-5.50062-6). URL: <https://doi.org/10.1016/B978-1-55860-332-5.50062-6> (cit. on pp. 294–296).
- [4] Andrei N. Tikhonov. “Systems of differential equations containing a small parameter multiplying the derivative”. In: *Matematicheskii Sbornik* 31(73).3 (1952). In Russian, pp. 575–586 (cit. on p. 779).
- [5] Hassan K. Khalil. *Nonlinear Systems*. 3rd. Upper Saddle River, NJ: Prentice-Hall, 2002 (cit. on pp. 779, 782).
- [6] Ali Saberi and Hassan K. Khalil. “Quadratic-type Lyapunov functions for singularly perturbed systems”. In: *IEEE Transactions on Automatic Control* 29.6 (1984), pp. 542–550. DOI: [10.1109/TAC.1984.1103586](https://doi.org/10.1109/TAC.1984.1103586) (cit. on p. 779).
- [7] P. Kokotovic, H. Khalil, and J. O’Reilly. *Singular perturbation methods in control: analysis and design*. 1986. DOI: [10.1137/1.9781611971118](https://doi.org/10.1137/1.9781611971118). URL: <https://doi.org/10.1137/1.9781611971118> (cit. on p. 779).
- [8] Lin-Yuan Chen, Nigel Goldenfeld, and Yoshitsugu Oono. “Renormalization group and singular perturbations: Multiple scales, boundary layers, and reductive perturbation theory”. In: *Physical Review E* 54.1 (1996), pp. 376–394. DOI: [10.1103/PhysRevE.54.376](https://doi.org/10.1103/PhysRevE.54.376) (cit. on p. 783).
- [9] Hermann Haken. *Synergetics: An Introduction. Nonequilibrium Phase Transitions and Self-Organization in Physics, Chemistry, and Biology*. 3rd ed. Berlin: Springer, 1983 (cit. on p. 783).
- [10] Karl Friston et al. “From pixels to planning: scale-free active inference”. In: *Frontiers in Network Physiology* 5 (2025), p. 1521963. DOI: [10.3389/fnetp.2025.1521963](https://doi.org/10.3389/fnetp.2025.1521963) (cit. on pp. 783, 784).